

Nullexit: Unified Project Document

Omar Alaaeldein

July 18, 2026

Contents

1	Project Directory Tree	4
2	README.md	6
3	agent.md	14
4	devref.md	22
5	diagrams.md	73
6	toggle.sh	78
7	recover.sh	106
8	setup.sh	115
9	scripts/setup-linux.sh	116
10	docker-compose.yml	117
11	scripts/common.sh	120
12	scripts/watcher.sh	135
13	scripts/crypto.sh	138
14	.aiignore	139
15	.claude/scheduled_tasks.lock	140
16	.claude/settings.local.json	140
17	.env.example	140
18	.gitignore	141
19	Recover-Gateway.applescript	142
20	Sweep-Gateway.applescript	142
21	Toggle-Gateway.applescript	142
22	benchmarks/benchmark.sh	143

23	benchmarks/test_load.py	144
24	black_list.txt	145
25	cloud-init.yaml	146
26	docker/tor/Dockerfile	146
27	docker/tor/entrypoint.sh	147
28	launchd/com.nullexit.noise.plist	147
29	launchd/com.nullexit.wake-recovery.plist	148
30	post-rules.txt	149
31	scripts/Dockerfile.routing-fix	150
32	scripts/Dockerfile.rule-compiler	150
33	scripts/Toggle-Gateway.bat	150
34	scripts/check_circuits.py	153
35	scripts/device-scramble.sh	154
36	scripts/diagnose-host-leak.sh	156
37	scripts/dns-proxy.py	165
38	scripts/dns_anomaly_detector.py	166
39	scripts/fix-docker-bridge-collision.sh	172
40	scripts/fix-ssh-delay.sh	177
41	scripts/generate_tex.py	178
42	scripts/logger/go.mod	181
43	scripts/logger/main.go	181
44	scripts/noise.sh	184
45	scripts/pf.conf	188
46	scripts/pihole-mode.sh	189
47	scripts/profiles.sh	191
48	scripts/reinstall-host-tailscale.sh	193
49	scripts/routing-fix.sh	202
50	scripts/rule-compiler/go.mod	206
51	scripts/rule-compiler/main.go	206
52	scripts/setup-common.sh	217

53	scripts/sweep.sh	223
54	scripts/unlock-files.sh	229
55	tor-bridges.txt	230
56	white_list.txt	230

1 Project Directory Tree

```
nullexit/
|-- .aignore
|-- .claude
|   |-- scheduled_tasks.lock
|   '-- settings.local.json
|-- .env.example
|-- .gitignore
|-- LICENSE
|-- README.md
|-- Recover-Gateway.applescript
|-- Sweep-Gateway.applescript
|-- Toggle-Gateway.applescript
|-- agent.md
|-- benchmarks
|   |-- benchmark.sh
|   '-- test_load.py
|-- black_list.txt
|-- cloud-init.yaml
|-- devref.md
|-- diagrams.md
|-- docker
|   '-- tor
|       |-- Dockerfile
|       '-- entrypoint.sh
|-- docker-compose.yml
|-- launchd
|   |-- com.nullexit.noise.plist
|   '-- com.nullexit.wake-recovery.plist
|-- post-rules.txt
|-- recover.sh
|-- scripts
|   |-- Dockerfile.routing-fix
|   |-- Dockerfile.rule-compiler
|   |-- Toggle-Gateway.bat
|   |-- check_circuits.py
|   |-- common.sh
|   |-- crypto.sh
|   |-- device-scramble.sh
|   |-- diagnose-host-leak.sh
|   |-- dns-proxy.py
|   |-- dns_anomaly_detector.py
|   |-- fix-docker-bridge-collision.sh
|   |-- fix-ssh-delay.sh
|   |-- generate_tex.py
|   |-- logger
|   |   |-- go.mod
|   |   '-- main.go
|   |-- noise.sh
|   |-- pf.conf
|   |-- pi-hole-mode.sh
|   |-- profiles.sh
|   '-- reinstall-host-tailscale.sh
```

```
| |-- routing-fix.sh
| |-- rule-compiler
| | |-- go.mod
| | |-- main.go
| |-- setup-common.sh
| |-- setup-linux.sh
| |-- sweep.sh
| |-- unlock-files.sh
| |-- watcher.sh
|-- setup.sh
|-- sweep-20260718T155650Z.txt
|-- sweep-20260718T160000Z.txt
|-- toggle.sh
|-- tor-bridges.txt
|-- white_list.txt
```

2 README.md

```
1 # nullexit: Tailscale + Cloudflare WARP Docker Gateway
2
3 > **Last updated:** July 12, 2026 [Architecture & Flow Diagrams ](./diagrams.md) [Full Dev Reference
4 ](./devref.md)
5
6 **nullexit** is a chained network gateway that routes all Tailscale exit-node traffic through a
7 Cloudflare WARP VPN tunnel double-encrypting every packet, hiding your ISP metadata, and providing
8 network-wide DNS ad-blocking (AdGuard Home) and kernel-level IP threat blocking ('ipset'/'iptables')
9 for every device on your mesh.
10
11 ---
12
13 ## 1. Prerequisites
14 - Docker and Docker Compose installed.
15 - A Tailscale account.
16 - **macOS:** Install the standalone Tailscale CLI via Homebrew ('brew install tailscale') and register '
17 tailscaled' as a system service ('brew services start tailscale'). No '.app' GUI install required.
18 - **OS & Python:** Developed and tested against **macOS 26.5.2**. The project has zero external Python
19 dependencies (no 'requirements.txt') and explicitly targets the default Apple-provided **Python
20 3.9.6** ('/usr/bin/python3'). While the scripts could function on newer versions, 'nullexit'
21 explicitly avoids experimenting with or modifying the built-in system Python environment to prevent
22 unexpected OS-level side effects.
23
24 ## 2. Installation & Setup
25
26 Run the interactive setup script for your OS. It downloads 'wgcf', generates fresh Cloudflare WARP
27 WireGuard keys, prompts for a Tailscale Auth Key, and writes your '.env'.
28
29 ```bash
30 ./setup.sh # macOS
31 ./scripts/setup-linux.sh # Linux
32 ```
33
34 ### Reference '.env'
35 This is a quick reference of the common variables. For the **complete, authoritative list** of every '.
36 env' variable (including all optional/advanced settings), see 'devref.md 7'.
37
38 ```env
39 WIREGUARD_PRIVATE_KEY=<Generated>
40 WIREGUARD_PUBLIC_KEY=<Generated>
41 WIREGUARD_ADDRESSES=<Generated>
42 TS_AUTHKEY=tskey-auth-...
43 GATEWAY_RULE_PROFILE=medium # light | medium | heavy
44 # Safe MSS for double-tunneled traffic (Tailscale + WARP). Don't raise this
45 # devref.md 15.3.2 found 1180 (and even 1160) still too large and causes the
46 # exact mysterious mid-transfer stalls this value exists to prevent.
47 GATEWAY_MSS=1120
48 # Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own
49 internet is NOT hijacked).
50 GATEWAY_HIJACK_HOST=true
51 GATEWAY_USE_EXIT_NODE=true
52 WARP_FAIL_THRESHOLD=6 # consecutive warp=off polls before auto-shutdown (default 6 = 30s)
53 KILL_SWITCH=true # enforce strict PF lock that breaks SSH if VPN fails. Leave true
54 # this is the core leak guarantee.
55 # On campus/enterprise networks (WPA2-Enterprise / 802.1x), AP Client Isolation blocks direct
56 # LAN traffic between devices. Tailscale attempts a local P2P upgrade anyway, which causes
57 # macOS gvproxy to SNAT-mangle the reply's source port poisoning the phone's WireGuard
58 # endpoint and blackholing 100% of its traffic. Setting this to false drops those local
59 # hole-punch packets so the phone safely falls back to Tailscale DERP relays.
60 # watcher.sh auto-detects 802.1x and AP isolation on every Wi-Fi change and overrides
61 # this setting automatically you only need to set this manually if auto-detection fails.
62 TAILSCALE_ALLOW_LAN_P2P=false # auto-managed; true only on trusted hotspots/home networks
63 ADGUARD_USER=admin # Username for AdGuard Home (defaults to admin if not set)
64 ADGUARD_PASSWORD=nullexit # Shared password for AdGuard Home and Tor ControlPort (defaults to
65 nullexit if not set)
66 BLOCKED_COUNTRIES="kp il" # dynamically block country IP ranges via ipdeny.com
67 TOR_EXCLUDE_EXIT_NODES={us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il} # comma-separated country codes to
68 exclude as Tor exit nodes
69 DEBUG_TRACE=false # true = mirror a full command-by-command trace of the toggle to
70 output.log (see 9)
71 ```
72
73 ## 3. Deploy the Gateway
74
75 **macOS (Recommended):**
76
77 ```bash
78 ./toggle.sh
79 ```
```

```

62 Running it again stops the gateway and restores your network. Handles the full lifecycle: Colima VM,
    containers, DNS hijacking, Tailscale exit-node routing, rule compilation, and sleep prevention.
63
64 **Linux / Manual:**
65 '''bash
66 docker compose up -d
67 '''
68
69 ## 4. Post-Deployment
70
71 ### Approve the Exit Node
72 1. Go to [Tailscale Admin Machines](https://login.tailscale.com/admin/machines).
73 2. Find the 'tailscale' device '...' **Edit route settings** enable as Exit Node.
74
75 ### Verify Traffic Flow
76 On any Tailscale client device, select this gateway as the exit node, then visit [api.ipify.org](https://
    api.ipify.org) or run:
77 '''bash
78 curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
79 # expect: warp=on
80 '''
81
82 ### AdGuard DNS Setup
83
84 To ensure client devices correctly route DNS through the AdGuard container (and do not bypass it via '
    tailscaled's internal resolver), you **must** configure your Tailscale Admin Console:
85 1. Go to the **DNS** tab in the [Tailscale Admin Console](https://login.tailscale.com/admin/dns).
86 2. Enable **MagicDNS**.
87 3. Under **Nameservers**, add a **Custom** nameserver and enter the gateway's Tailscale IPv4 address (run
    'tailscale ip -4' on the gateway machine to find it, e.g., '100.x.x.x').
88 4. Click the **three dots (...)** next to the nameserver you just added and enable **"Use with exit node
    "**.
89 5. Delete any other Global Nameservers (like Google or Cloudflare).
90 6. Toggle ON **Override local DNS**.
91
92 Traffic will now be correctly forced into the AdGuard sinkhole, and mobile devices will be blocked from
    bypassing it via encrypted DNS (DoH/DoT).
93
94 > [!WARNING]
95 > **Disable "Private DNS" on mobile devices.** Android's Private DNS (DoT/DoH) bypasses port 53
    interception entirely. Go to 'Settings Connections More connection settings Private DNS Off'.
96
97 > [!WARNING]
98 > **Bandwidth doubling:** Streaming 1GB of video on a remote device consumes 1GB download + 1GB upload on
    the host's network. Be mindful on metered connections.
99
100 ### AdGuard Dashboard
101 Navigate to 'http://<gateway-tailscale-ip>:3000' credentials: **'admin' / 'nullexit'**.
102
103 ### Access Any Mesh Device From Anywhere (SFTP/SSH)
104
105 Once the gateway exit node is online, **any device on your Tailscale mesh can reach any other mesh device
    , anywhere in the world** as long as both devices are on the mesh and the target device has an SSH
    server running. No port forwarding, no public IPs, no firewall holes. You can SSH into any device
    using its MagicDNS hostname (e.g. 'ssh username@my-macbook') or its assigned Tailscale IP ('100.x.x.
    x').
106
107 > **Security Tip (Zero Local Attack Surface):** It is highly recommended to disable macOS's native "
    Remote Login" and "Screen Sharing" in System Settings to prevent exposing those ports (22, 5900) to
    physical Wi-Fi networks (which scanners can fingerprint). The gateway enforces 'tailscale up --ssh=
    true' and the 'pf' Kill-Switch automatically blocks local Wi-Fi access to these ports. Since both
    MagicDNS and Tailscale IPs route through the exact same encrypted tunnel, they are equally securebut
    **MagicDNS is recommended simply for ease of use.**
108
109 ### Extreme OPSEC: The Secret Tor-over-WARP Proxy
110 'nullexit' includes a completely isolated, headless Tor infrastructure proxy designed for terminals and
    hacking tools (e.g., 'curl', 'sqlmap', 'nmap').
111 To protect against local malware fingerprinting:
112 1. **Procedural Ports:** The proxy ports are randomized into the ephemeral range (10000-65000) on every
    single boot using a strict kernel socket collision-avoidance check ('netstat').
113 2. **Honey-Port Tripwire:** The gateway spawns a fake trap port. If any malware attempts a sequential
    port-scan on 'localhost' to find the proxy, it trips the Honey-Port, instantly triggering a macOS
    desktop notification and a critical log alert.
114
115 > [!NOTE]
116 > **Threat Model limitation:** Port randomization and the Honey-Port only defeat blind, generic port
    scanners. They do not protect against targeted local malware that runs 'lsof -i' or reads the '/tmp/
    nullexit-ports.env' file directly. To actually prevent malicious local processes from reaching the
    proxy, you must run a host-level outbound firewall (like LuLu 4.3.0+ or Little Snitch) with **
    localhost/loopback filtering explicitly enabled** to gate access by process identity.
117

```

```

118 To use the Tor proxy, run 'cat /tmp/nullexit-ports.env' to find your '$TOR SOCKS_PORT' for the current
    session, and point your hardened browser (LibreWolf/Tor Browser) or CLI tool to '127.0.0.1:
    $TOR SOCKS_PORT'.
119
120 > [!WARNING]
121 > **Preventing DNS Leaks (SOCKS5 vs. SOCKS5-Hostname):**
122 > When routing command-line tools through Tor, you must force the tool to delegate DNS resolution to
    the Tor proxy. Using plain SOCKS5 (e.g., 'socks5://' or curl's '--socks5' flag) will resolve
    hostnames locally using the host's DNS settings (AdGuard/WARP) before establishing the connection.
    While this DNS traffic is encrypted inside the WARP tunnel, the destination domain name is still
    leaked to AdGuard and Cloudflare, undermining your anonymity.
123
124 > Use these correct protocols and configurations:
125 > - 'curl': Use '--socks5-hostname' or the 'socks5h://' scheme (e.g., 'curl -x socks5h://127.0.0.1:
    $TOR SOCKS_PORT https://check.torproject.org'). Do not use '--socks5' or 'socks5://'.
126 > - 'sqlmap': Pass the proxy URL using the 'socks5h://' scheme to ensure remote resolution: '--proxy
    =socks5h://127.0.0.1:$TOR SOCKS_PORT'.
127 > - 'nmap': Nmap's built-in '--proxies' option resolves hostnames locally. To safely scan targets
    without DNS leaks, tunnel it using 'proxychains-ng' configured with 'proxy_dns' enabled (add 'socks5
    127.0.0.1 <PORT>' to '/etc/proxychains.conf' or a local config, then run 'proxychains4 nmap <args
    >').
128
129
130 ##### Transparent .onion Browsing & Dark Web Access
131 In addition to the randomized SOCKS5 proxy port, 'nullexit' provides seamless, transparent '.onion'
    domain browsing for all devices on your Tailscale mesh:
132 - Zero-Config DNS: AdGuard Home automatically intercepts queries for '.onion' domains and resolves
    them to Tor's virtual mapping subnet.
133 - Tailscale Auto-Routing: The virtual subnet is mapped to the '198.18.0.0/15' benchmark range (RFC
    2544). Since this range is not a private IP subnet, Tailscale's Exit Node mode automatically routes
    it to the gateway (bypassing local LAN exclusions).
134 - Transparent NAT Proxying: Kernel firewall rules in the gateway redirect all TCP traffic destined
    for '198.18.0.0/15' directly into Tor's transparent port ('9040').
135 - Exit Node Exclusion: Supports the ability to exclude specific countries from being used as Tor exit
    nodes via the 'TOR_EXCLUDE_EXIT_NODES' environment variable in '.env'.
136
137 ##### Enable in Web Browsers
138 To prevent accidental DNS leaks, web browsers block '.onion' lookups by default. To browse dark web sites
    natively:
139 - Firefox: Go to 'about:config', search for 'network.dns.blockDotOnion', and set it to 'false'. Type
    any onion address (e.g. 'duckduckgogg42xjoc72x3sjasowarfbgcmvfimaftt6twagswzczad.onion') directly
    into the URL bar and it will load.
140 - Mobile Browsers (iOS/Android): Connect to your Tailscale exit node, open Firefox (with
    blockDotOnion disabled), and browse onion sites natively on the go.
141
142 ##### Hardening: Using obfs4 Tor Bridges
143 If you want to disguise your Tor traffic from the WARP exit node provider, you can optionally enable Tor
    bridges:
144 1. Obtain private 'obfs4' bridge lines from [bridges.torproject.org](https://bridges.torproject.org) or
    the official Telegram bot.
145 2. Create a file named 'tor-bridges.txt' in the root of the 'nullexit' folder and paste your bridge lines
    inside it (one per line).
146 3. Set 'TOR_USE_BRIDGES=true' in your '.env' file and restart the gateway. The proxy will refuse to start
    if the bridges file is empty or missing.
147
148 ##### Example: Browse your Android phone's files from your Mac
149
150 1. On your Android phone, install [Termux](https://termux.dev/) and [Termux:Boot](https://f-droid.org
    /packages/com.termux.boot/) from F-Droid.
151 2. In Termux, install and start the SSH server:
152   ``bash
153   pkg install openssh
154   sshd
155   ``
156   Termux defaults to port 8022 (ports below 1024 require root on Android).
157 3. Find your phone's Tailscale IP:
158   ``bash
159   tailscale ip -4 # e.g. 100.87.42.87
160   ``
161 4. Stop Android from killing Termux do ALL of these (Samsung One UI is especially aggressive):
162   - Settings Apps Termux Battery Unrestricted
163   - Device Care Battery Background usage limits Sleeping apps / Deep sleeping apps remove Termux
    if listed
164   - In Termux, run 'termux-wake-lock' to hold a persistent wakelock
165 5. On your Mac/PC, open [Cyberduck](https://cyberduck.io/) (or any SFTP client: WinSCP, FileZilla, FE
    File Explorer on iOS) and connect to:
166   - Server: '100.87.42.87' (your phone's Tailscale IP)
167   - Port: '8022'
168   - Protocol: SFTP (SSH File Transfer Protocol)
169   - Username: (your Termux username, usually 'u0_a123' run 'whoami' in Termux to check)
170

```



```

171 - **Password:** (your Termux password set with 'passwd' in Termux if you haven't already)
172
173 That's it you can now browse, download, and upload files to your phone from any device on the mesh, no
174 matter where either device is physically located. This works identically for any mesh device: Mac
175 Mac, Windows Linux, phone NAS, etc.
176
177 ---
178
179 ## 5. Multi-Layer Firewall & Kill-Switch
180
181 ### Layer 1 DNS Sinkhole (AdGuard Home)
182 Blocks ads and tracking domains before they are downloaded. In automated Lighthouse tests on ad-heavy
183 sites:
184 - **Time to Interactive:** ~22% faster (51.0s 41.8s)
185 - **Speed Index:** ~20% faster (15.3s 12.8s)
186
187 Customize blocking via 'black_list.txt' and 'white_list.txt'. Rule profiles ('light' / 'medium' / 'heavy
188 ') are set in '.env'.
189
190 ### Layer 2 Kernel IP Firewall ('ipset'/'iptables')
191 On every startup, the 'rule-compiler' fetches threat-intelligence feeds (Spamhaus DROP, Feodo Tracker,
192 Emerging Threats, CINS) and compiles **~16,700 unique malicious IPs/CIDRs** into the kernel 'FORWARD
193 ' chain blocking both outbound C2 connections and inbound attack traffic. Zero configuration
194 required.
195
196 ### Layer 3 Geo-IP Blocking
197 Dynamically blocks all traffic to and from specific countries using live IP ranges from 'ipdeny.com'.
198 To add countries to your blocklist, open your '.env' file and set the 'BLOCKED_COUNTRIES' variable with
199 2-letter ISO country codes (e.g. 'BLOCKED_COUNTRIES="kp il cn ru"'), then restart the gateway.
200
201 ### Layer 4 macOS PF Kill-Switch
202 A native macOS Packet Filter ('pf') kill-switch that enforces a strict default-deny policy on your
203 physical Wi-Fi interface ('en0'). When 'KILL_SWITCH=true' is set in your '.env', it drops all
204 outgoing traffic except the Cloudflare WARP endpoints and Tailscale DERP relays.
205
206 - **Failsafe Design:** If the VPN tunnel crashes or Docker dies, your Mac completely loses internet
207 access rather than leaking your real IP.
208 - **Remote-Access Safe:** Carefully designed to whitelist Tailscale's control plane ('192.200.0.0/24'), '
209 utun*' tunnel traffic, and local LAN subnets. This guarantees that your remote SSH sessions and
210 local AirDrop will survive even when the kill switch engages.
211 - **Setup Requirement:** The toggle scripts need permission to manipulate the network and firewall in the
212 background without prompting you for a password. You must configure your passwordless sudo config
213 by running exactly:
214
215 ```bash
216 echo "$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /usr/bin/dscacheutil, /usr/bin/
217 killall, /usr/bin/pkill, /bin/kill, /sbin/route, /sbin/ifconfig, /usr/bin/true, /opt/homebrew/bin/
218 brew, /usr/local/bin/brew, /usr/bin/python3 -I $PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3
219 -I $PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I $PWD/scripts/dns-proxy.py" | sudo tee /etc/
220 sudoers.d/nullexit
221 ```
222
223 ---
224
225 ## 6. Toggle Scripts
226
227 ### macOS 'toggle.sh'
228 Fully automated lifecycle script. Also supports a native '.app' launcher:
229
230 ```bash
231 osascript -o "Toggle Gateway.app" Toggle-Gateway.applescript
232 ```
233
234 Optional '.env' settings (common subset; the full canonical table lives in 'devref.md 7'):
235 - 'GATEWAY_BYPASS_PING=true' proceed even if pre-flight connectivity checks fail.
236 - 'GATEWAY_USE_EXIT_NODE=false' DNS-only mode, skip exit-node routing.
237 - 'GATEWAY_HIJACK_HOST=false' skip DNS hijacking on the host (provides VPN/adblocking only to external
238 Tailscale peers).
239 - 'GATEWAY_MSS' override the TCP MSS clamp for the double-tunnel. Default (and only validated-safe value
240 ) is '1120'; devref.md 15.3.2 found 1180 and even 1160 still too large for the Tailscale+WARP
241 double-encapsulation overhead, causing the exact mid-transfer stalls this clamp exists to prevent.
242 Applies to both the container ('post-rules.txt') and the host ('pf.conf') they're kept in sync
243 automatically.
244 - 'HOST_MTU=1200' override the host Tailscale interface MTU. Defaults to 1200 to fit inside the 1280-
245 byte WARP tunnel.
246 - 'KILL_SWITCH=true' enforce strict network lock that breaks SSH if the VPN fails.
247 - 'STOP_COLIMA_ON_EXIT=true' fully shut down the Colima VM on toggle-off (saves battery on dedicated
248 hosts).

```

```

223 - 'TAILSCALE_ALLOW_LAN_P2P=true' allow direct Tailscale LAN P2P connections between devices on the same
224 network.
225 - **Default: 'false'.** On campus or enterprise Wi-Fi (WPA2-Enterprise / 802.1x), AP Client Isolation
blocks local device-to-device traffic. When Tailscale still attempts a local P2P upgrade, macOS's '
gvproxy' layer SNAT-mangles the reply's source port. WireGuard's endpoint roaming treats this as a
legitimate address change and updates the phone's outbound path to the newly-mangled port which has
no listener. This blackholes 100% of the phone's traffic while the DERP relay path is abandoned.
Setting this to 'false' preemptively drops all RFC1918-destined Tailscale UDP from the gateway so
the phone receives a clean failure and safely falls back to DERP.
226 - **Auto-managed:** 'watcher.sh' automatically detects WPA2-Enterprise via 'system_profiler
SPAirPortDataType' (the 'airport' binary was removed in macOS 14.4) and probes for AP isolation by
pinging the default gateway ('route -n get 1.1.1.1') on every network change. It writes the detected
value to '.lan_p2p_detected' in the repo root, which 'routing-fix.sh' re-reads every 30 seconds **
no restart required.** Only set this manually if auto-detection fails. See 'devref.md 15.7.1' for
the full SNAT endpoint poisoning analysis and 'devref.md 15.11.6' for the 'airport' removal
background. *(Note: Direct P2P between the gateway container and the host Mac itself is always
permitted via the Docker bridge to bypass DERP relays, regardless of this setting).*
227 - Set to 'true' on trusted home networks or Windows/mobile hotspots (no AP isolation) for a faster,
lower-latency direct path.
228 - 'WARP_FAIL_THRESHOLD=3' number of consecutive failed checks before forcing a gateway teardown (default
: 6 checks = 30s).
229 - 'WARP_ENDPOINT_1' / 'WARP_ENDPOINT_2' override the Cloudflare WARP WireGuard endpoints.
230
231 ### Linux 'toggle.sh' / 'recover.sh'
232 'bash
233 ./toggle.sh # toggle ON/OFF (cross-platform, dispatches on $OSTYPE)
234 ./recover.sh # recovery tool (cross-platform)
235 '
236
237 ### Sleep Prevention (macOS)
238 When the gateway is ON, 'caffeinate -i' runs in the background to prevent idle system sleep, keeping
239 Docker and all services alive even on battery. The display can still sleep normally.
240
241 ### Auto-Recovery (macOS post-wake / post-roam)
242 Install the launchd LaunchAgent so the gateway self-heals after sleep/wake and Wi-Fi roams:
243 'bash
244 cp ./launchd/com.nullexit.wake-recovery.plist ~/Library/LaunchAgents/
245 sed -i 's|_NULLEXIT_HOME_|$(pwd)|' ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
246 launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
247 '
248
249 Wi-Fi roams and network changes are caught reliably. Lid-close/wake is
250 **best-effort on macOS Sonoma** Apple suspends LaunchAgents during forced
251 sleep rather than re-triggering them on every wake, so the watcher can miss
252 the specific wake event (other self-healing the DNS Watcher, Tailscale's
253 DERP fallback, gluetun's healthcheck still recovers within ~30-60s
254 regardless). If lid-wake responsiveness matters more than that window,
255 'brew install sleepwatcher' and hook 'recover.sh --post-wake' to '~/.wake'
256 instead. See 'devref.md 15.5.1' for the full deep dive.
257
258 ### Config profiles taming the flag sprawl
259 The opt-in features below add up to several interacting '.env' switches. Rather than
260 hand-tune them, 'scripts/profiles.sh' bundles the intent-level flags into three
261 known-good presets. It is **config-only** it just edits '.env', never touches the
262 live gateway so changes apply on your next './toggle.sh --restart', and it forces
263 'KILL_SWITCH=true' in every profile.
264
265 'bash
266 bash scripts/profiles.sh show # the whole current config, grouped + explained
267 bash scripts/profiles.sh preview home # exactly what 'apply' would change (read-only)
268 bash scripts/profiles.sh apply home # write it (backs up .env; never restarts)
269 '
270
271 Profiles: **campus** (AP-isolated enterprise, conservative) **home** (trusted, fast
272 direct P2P, FaceTime + LAN Pi-hole) **hardened** (cover traffic + Tor bridges, no
273 real-IP exposure).
274
275 ### Cover traffic / dummy padding (macOS optional, opt-in)
276 'scripts/noise.sh' emits verifiable-random UDP padding toward the exit so a passive
277 on-path observer sees a continuously varying encrypted flow instead of your real
278 traffic envelope (a metadata / traffic-flow correlation defence). Unlike the rest
279 of the toolkit it is **active** it puts traffic on the wire so it stays inert
280 unless you opt in. **One switch controls everything:** set 'NOISE_ENABLED=true' in
281 '.env'. There is no separate "persistent vs one-shot" option installing the
282 LaunchAgent below simply means padding auto-starts whenever 'NOISE_ENABLED=true',
283 and the job is a clean no-op whenever it is 'false'.
284
285 'bash
286 # One-off (this session only):
287 bash scripts/noise.sh verify # prove the noise is statistically random (entropy/monobit/chi-square)
288 bash scripts/noise.sh start # arm padding now; 'status' / 'stop' to manage

```

```

287 # Persistent (survives reboot/login) install the LaunchAgent once:
288 cp ./launchd/com.nullexit.noise.plist ~/Library/LaunchAgents/
289 sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" ~/Library/LaunchAgents/com.nullexit.noise.plist
290 launchctl load -w ~/Library/LaunchAgents/com.nullexit.noise.plist
291 # To stop persisting: launchctl bootout gui/$UID/com.nullexit.noise
292 ''
293 Tunables (rate, packet size, jitter, target) live in '.env' under the
294 'NOISE_PAD_*' keys; defaults are a modest 64 kbps toward the gateway. Honest
295 limits: this is one-way random padding it blurs uplink volume + timing, not a
296 full traffic-analysis-resistant shaper. See the 'noise.sh' / 'Cover Traffic &
297 Padding' notes in the knowledge graph.
298
299 ### LAN Pi-hole mode (macOS optional, opt-in)
300 nullexit is already a Pi-hole for **mesh** devices AdGuard Home filters DNS for
301 everything that routes through the gateway over Tailscale. This mode extends that
302 filtering to **non-Tailscale** devices on the same LAN (a smart TV, a guest, IoT):
303 point their DNS server at this Mac's IP and they get the same ad/tracker/threat
304 blocking. It reuses 'dns-proxy.py' bound to the LAN interface, forwarding to
305 AdGuard.
306
307 ''bash
308 # .env: PIHOLE_LAN_MODE=true (default false)
309 sudo bash scripts/pihole-mode.sh enable # bind the LAN DNS forwarder (:53 needs root)
310 bash scripts/pihole-mode.sh test # prove it resolves clean + sinkholes blocked
311 bash scripts/pihole-mode.sh status # running? listen addr, AdGuard reachability
312 sudo bash scripts/pihole-mode.sh disable
313 ''
314
315 Honest limits: it only works on a **trusted, non-isolated** network on
316 AP-isolated WPA2-Enterprise (the same isolation that forces phones onto DERP)
317 other devices can't reach this Mac at all. And it's **DNS filtering only** it
318 does not route those devices through WARP. See the 'pihole-mode.sh' note in the
319 knowledge graph.
320
321 ### Device Scramble Mode (macOS optional, opt-in)
322 Randomizes the identifiers a network/observer uses to fingerprint your device (hostname and MAC address).
323 This is useful on open / no-auth Wi-Fi (caf, hotel, airport) where these IDs are the only handle on
324 you.
325 * Note: MAC spoofing requires root. The 'test-mac' command is fail-safe and detects Apple-Silicon spoof-
326 blocking without stranding you.
327
328 ''bash
329 bash scripts/device-scramble.sh status # show current + saved identity
330 sudo bash scripts/device-scramble.sh scramble # random hostname (add --mac to also rotate MAC)
331 sudo bash scripts/device-scramble.sh restore # put your real identity back
332 sudo bash scripts/device-scramble.sh test-mac # safe test to see if this network takes a random MAC
333 ''
334
335 ---
336
337 ## 7. Networking Notes
338
339 - **Port conflicts:** None. Tailscale uses dynamic UDP ports; WARP uses outbound UDP:2408. The gateway
340 binds '0.0.0.0:41642/udp' on the host for Tailscale's WireGuard listener this is intentional and
341 required to receive inbound hole-punch packets from peers on the internet. All other internal ports
342 (AdGuard :5335, SOCKS5 :1080) are bound to '127.0.0.1' only and are not reachable from the LAN.
343 - **Tailscale P2P vs DERP:** On the same Wi-Fi, Tailscale establishes a fast P2P connection. On cellular
344 (CGNAT), it always falls back to a DERP relay (~80-200ms). See 'devref.md 5.1'.
345 - **AirDrop / AirPlay:** Unaffected the gateway only touches standard Wi-Fi/Ethernet interfaces, not '
346 awd10'.
347 - **VPN coexistence:** Do **not** run a local VPN client (WARP, Mullvad, NordVPN) simultaneously with the
348 exit node. Both fight for the default route and will blackhole your connection.
349 - **Banking & financial sites:** May intermittently block logins through WARP. Banks use anti-fraud
350 databases that flag datacenter IP ranges (like Cloudflare's '104.28.x.x') as VPN/proxy traffic.
351 Success fluctuates depending on which WARP IP you land on. If a site blocks you, either disable the
352 exit node temporarily for that session or use the bank's mobile app (which often uses different
353 fraud detection).
354 - **MSS clamping (Bufferbloat):** Double-tunnelling adds ~120-160 bytes of overhead per packet. This is
355 now automatically handled natively via the macOS 'pf' firewall, which universally applies TCP MSS
356 Clamping to all outbound packets to strictly prevent MTU fragmentation and bufferbloat. No manual '.
357 env' configuration is required. Additionally, the host's Tailscale interface MTU is explicitly
358 lowered to 1200 (configurable via 'HOST_MTU' in '.env') to guarantee packets fit inside the
359 container's WARP tunnel without fragmenting.
360
361 ---
362
363 ## 8. Privacy & Security Hardening
364
365 The stack shifts trust across four layers: WARP hides traffic from your ISP, AdGuard blocks tracking at
366 DNS, Tailscale encrypts device-to-device, and 'ipset' blocks known malicious IPs at the kernel.

```

```

349 For higher security, the gateway supports these drop-in upgrades:
350
351 | Layer | Default | Hardened |
352 |---|---|---|
353 | VPN Exit | Cloudflare WARP (US) | Mullvad (Swedish, no-logs, anonymous payment) |
354 | DNS Resolver | AdGuard via WARP | AdGuard via Mullvad DoH |
355 | Coordination | Tailscale (US) | Headscale (self-hosted) |
356 | Auth | Persistent OAuth keys | Ephemeral auth keys |
357 | Content | WireGuard asymmetric | WireGuard + Pre-Shared Keys (PSK) |
358
359 ##### Python Runtime Airgap (Isolated Mode)
360 **Never** install third-party 'pip' packages (like 'requests' or 'pyyaml') into the Python environment
    that 'nullexit' uses. All python scripts in this project (e.g., the 'dns-proxy.py' daemon) are
    strictly engineered to run on the zero-dependency Python Standard Library. To strictly isolate the
    environment from supply-chain attacks, 'nullexit' executes these scripts using **Python's Isolated
    Mode** ('python3 -I'). This deliberately blinds the execution environment to any malicious user-
    installed 'pip' packages on the host, permanently air-gapping the gateway from PyPI.
361
362 See 'devref.md 9' for the full threat model, Snowden programme analysis, and trust assumptions.
363
364 ---
365
366 ## 9. Debugging
367
368 - **'output.log'** All stderr from 'toggle.sh', 'recover.sh', 'setup.sh', and the rule-compiler is
    written here. The WARP Watcher ('WARP_FAIL_THRESHOLD') also writes its events here 'WARP DOWN', '
    WARP RECOVERED', and 'WARP SHUTDOWN' notices. Check this first.
369 - Every run now records lifecycle breadcrumbs regardless of 'DEBUG_TRACE': an 'EXEC:' / 'EXIT <code>:'
    line brackets each heavyweight command ('colima start', 'docker compose up -d --build', 'docker
    compose down'), and an 'EXIT: code=<N> phase=<init|stop|start>' line is written when the script
    exits **for any reason** including a killed terminal, 'pkill', or OOM. This closes a prior blind
    spot where an interrupted **START** (e.g. a '--restart' race while Wi-Fi was re-associating) left
    the log ending abruptly after the teardown with no indication of what failed. If a toggle "does
    nothing" or the gateway ends up half-up, 'grep -E 'EXEC:|EXIT ' output.log | tail' shows the last
    command run and its result.
370 - **'DEBUG_TRACE=true'** (in '.env') Full command-by-command execution trace of 'toggle.sh', mirrored to
    'output.log' for deep post-mortems of intermittent failures. Each traced line is timestamped and
    tagged with 'file:line:function' (via a custom 'PS4'), so you can follow the *exact* code path a run
    took including subshells, Docker/Colima calls, and where it aborted. Off by default because it is
    verbose (the log grows quickly); enable it only while reproducing a specific problem, then set it
    back to 'false'. Implementation note: it prefers bash's 'BASH_XTRACEFD' (bash 4.1, e.g. most Linux)
    to send the trace to the log while keeping your terminal clean; on macOS's stock '/bin/bash' 3.2 it
    transparently falls back to tee'ing stderr to the log (trace also appears in the terminal). It never
    uses the dynamic-fd ('exec {var}>>') form, so it is safe on 3.2. To read a trace: 'grep -nE '^\+ '
    output.log'.
371 - **'bash scripts/diagnose-host-leak.sh'** One-shot host-routing diagnostic. Runs 8 checks and
    classifies your state into one of four known scenarios (System Extension conflict, SOCKS5 fallback,
    IPv6 leak, route-table freeze) or OK, and prints the exact fix command. Check 5 uses 'route -n get
    1.1.1.1' to ask the kernel which interface real traffic actually resolves through (more accurate
    than 'netstat -rn | grep default' on macOS). Pass '--fix' to apply the remediation automatically.
    Pass '--watch' (or '--watch 30') to run a full baseline diagnostic then continuously monitor warp/
    IPv6/default-route for leaks, alerting on any state change.
372 - **'bash scripts/sweep.sh'** One-shot, read-only gateway health sweep. Runs 7 checks containers up/
    healthy, double-tunnel/IP-leak (host egress == container egress with 'warp=on'), direct P2P to the
    gateway (no DERP relay), AdGuard compiled-rule counts + live DNS interception, PF kill-switch armed,
    tailnet reachability (a 'tailscale ping' to every online peer, flagging DERP-relayed peers as
    throughput-degraded), and real throughput (host path vs. a WARP-only baseline via the container's
    SOCKS5 proxy, against a fast CDN no exec/install inside the container) then rolls them into a PASS
    /WARN/FAIL verdict. Writes a timestamped 'sweep-<UTC>.txt' report to the repo root (diffable across
    runs) and exits non-zero on FAIL, so it can gate launchd hooks or CI. Pass '--quiet' for just the
    verdict line. It never mutates routing, PF, DNS, or containers safe to run against a live gateway
    at any time.
373 - **'bash scripts/fix-docker-bridge-collision.sh'** One-shot fix for Docker bridge IP conflicts. If your
    local Wi-Fi router assigns you an IP in the '172.17.0.0/16' range, it will violently collide with
    Docker's default internal bridge, permanently killing your internet. This script detects the
    collision and auto-patches your Colima VM to use a safe subnet (e.g. '10.200.0.1/24').
374 - **'devref.md'** Complete architecture reference, routing deep-dives, and full resolved-issues log.
    Read this before modifying any routing or firewall logic.
375
376 > [!CAUTION]
377 > **Two infinite routing loops are possible.** (1) If a hotspot device providing internet to the gateway
    host is *also* using the gateway as its exit node, packets bounce forever between the two. (2) When
    the host's default route is redirected to 'utun*', WARP endpoint packets can loop back into the
    tunnel. Both are documented in 'devref.md 15.2.1' and '15.2.2' and are automatically handled by '
    toggle.sh'.
378
379 ---
380
381 ## 10. Unified Project Document (Quine)
382

```

```

383 The project includes a 'generate_tex.py' script that compiles the entire source code, configurations, and
384 documentation into a single, unified LaTeX document ('nullexit_unified.tex' / '.pdf').
385
386 Funnily enough, this document acts as a "project-level quine." It encapsulates the entirety of its own
387 source code, its documentation, and the exact Python code required to generate itself. It serves as
388 a self-contained, long-term archiving strategy providing everything needed to reconstruct and
389 understand the 'nullexit' system from a single file.
390
391 ---
392
393 ## 11. Acknowledgements
394 - **[SyameimaruKoa](https://github.com/SyameimaruKoa):** Dual-stack TCP MSS clamping, 'SIGHUP' state-
395 tracking in the routing sidecar, and 'TS_AUTH_ONCE' integration.
396
397 ## 12. License
398 GNU Affero General Public License v3. See [LICENSE](./LICENSE).
399
400 ## 13. Changelog
401
402 - **July 18, 2026** Fixed two real bugs surfaced by live sweep evidence: the QUIC/DoT 'REJECT' rules in
403 'post-rules.txt' were shadowed by an earlier blanket 'ACCEPT' (0 packets ever matched the intended
404 Instagram/QUIC-fragmentation fix from 'devref.md 15.3.3' was never actually enforced), now reordered
405 ahead of it; and the host-side 'pf.conf' MSS clamp was stuck at the pre-'15.3.2' value of '1160'
406 while the container had already moved to '1120' 'scripts/common.sh' now syncs both from '
407 GATEWAY_MSS' so they can't drift apart again. Also fixed 'setup.sh'/'scripts/setup-common.sh' never
408 generating 'NULLEXIT_SEED' (every fresh install hard-failed its first 'toggle.sh' on a missing-seed
409 crypto error) and defaulting 'KILL_SWITCH' to 'false' for new installs see 'devref.md 15.12.23'. '
410 scripts/sweep.sh' gained a 7th check (**throughput**, host path vs. a WARP-only baseline against a
411 fast CDN) see 'devref.md 15.3.6' for why the target choice matters and '9' below for the check
412 itself.
413
414 - **July 17, 2026** 'scripts/sweep.sh' (the read-only automated health sweep added July 15, mirroring '
415 sweep.md 1') gained a 6th check: **tailnet reachability** every Tailscale peer reported online is
416 probed with 'tailscale ping' (TSMP, so it works even on phones that drop ICMP) and rolled into the
417 PASS/WARN/FAIL verdict. Documented in 9 above.
418
419 - **July 12, 2026** Tor ControlPort dynamic password authentication (unified with 'ADGUARD_PASSWORD'),
420 enabling live circuit inspection via '/sweep'. See 'devref.md 15.10.2'.
421
422 - **July 11, 2026** Tor container hardening: 'obfs4proxy' restored (base image 'debian:bookworm-slim'),
423 robust bridge-file validation, and image pinning. See 'devref.md 15.10.1'.
424
425 - **July 10, 2026** Roaming/startup stabilization suite: fixed the Wi-Fi-roam host-IP leak (raw ISP
426 '132.x') caused by the WARP Watcher racing post-wake recovery, plus a batch of pre-flight/
427 concurrency/kill-switch/Tailscale-flag bugs and the macOS SNAT endpoint-poisoning P2P blackhole.
428 Full post-mortems in the Incident Log see 'devref.md 15.2.7', '15.5.3', '15.5.5', '15.2.8', '15.2.9',
429 '15.7.1', '15.7.3', '15.7.5', '15.12.18'.
430
431 - **July 6, 2026** Edge-case security and stability hardening: DNS Watcher interface independence,
432 robust lock-file PID verification, fail-closed crypto integrity check, strict 'iptables' pinning for
433 tailscale0tun0, Docker port binding to '127.0.0.1' to prevent LAN exposure, routing verification
434 via 'netstat -nr'. See 'devref.md 17'.
435
436 - **July 4, 2026** Colima bridged networking, SOCKS5 proxy migrated natively to Tailscale, headless
437 prompt crashes fixed, proxy bypass domains added to fix LAN P2P. Python dependencies completely
438 eliminated: 'logger' and 'rule-compiler' ported to blazing-fast Go multi-stage Docker builds. Added
439 'crypto.sh' to enforce cryptographic signature validation on all core bash scripts. Massively
440 refactored 'toggle.sh' and 'recover.sh' to eliminate 200 lines of duplicated code by migrating
441 network and process logic to 'scripts/common.sh'. Added a concurrency mutex to 'toggle.sh', resolved
442 unbound 'TS_AUTHKEY' check crashes on setup, and integrated Go validation to the CI pipeline.
443
444 - **July 3, 2026** **Critical Security Updates:** Addressed the overnight silent IP leak and improved
445 geo-blocking. Introduced 'scripts/host-leak-probe.sh' (a sub-second host-egress leak detector) and a
446 statistical threshold auto-shutdown watcher. Added 'unlock-files.sh' to safely reset file
447 permissions. Resolved numerous startup regressions and system bugs.
448
449 - **July 2, 2026** Two infinite routing loops resolved ('devref.md 15.2.2'): host-VM tunnel loop (WARP
450 bypass routes now via gateway IP + manual 'utun*' redirection) and Docker Compose subnet takeover
451 ('10.200.1.0/24' lock). '--accept-routes=true' now explicit on all 'tailscale up --reset' calls.
452
453 - **July 2, 2026** Auto-recovery daemon documented in 6 above ('scripts/watcher.sh' + launchd plist).
454 See 'devref.md 15.5.1'.
455
456 - **July 2, 2026** Linux scripts ('scripts/toggle-linux.sh', 'recover-linux.sh', 'setup-linux.sh')
457 documented in 6.
458
459 - **July 1, 2026** 'recover.sh --post-wake' ships; wired to watcher daemon. Triggered on every sleep/
460 wake or network-state change.
461
462 - **June 28, 2026** 8 internal scripts moved to 'scripts/'. User-facing files ('toggle.sh', 'recover.sh'
463 ', 'setup.sh', 'docker-compose.yml') stay at repo root.
464
465 - **June 28, 2026** All hardcoded install paths removed. 'SCRIPT_DIR'--relative resolution and '
466 _NULLEXIT_HOME_' placeholder used throughout.
467
468 ### Cryptographic Integrity Verification
469
470 nullexit uses a built-in cryptographic integrity checker designed to provide tamper-evidence* against
471 accidental edits, logic bugs, or non-privileged malware. During setup, a 256-bit 'NULLEXIT_SEED' is
472 injected into your '.env' file. The core entrypoint scripts ('toggle.sh', 'recover.sh', etc.) are
473 hashed using HMAC-SHA256, and their signatures are stored in '.signatures'.
474
475 *(Note: This is not a strict boundary against an attacker who already has write access to the repo, as
476 they could simply read the seed and re-sign the scripts. It is designed as a strict footgun-catcher)

```

```

.*
417
418 If any script is modified without authorization, the gateway will loudly fail to boot. If you make
    intentional edits to the bash scripts, you must re-sign them by running:
419 ``bash
420 ./scripts/crypto.sh --sign
421 ``

```

3 agent.md

```

1 # nullexit Detailed Agent Analysis
2
3 > Complete per-file breakdown of every function, constant, duplication, and bug found.
4 > **Note (July 2026):** 'sync-rules.py' and 'logger.py' have been fully ported to Go ('scripts/rule-
    compiler/main.go' and 'scripts/logger/main.go') to dramatically improve performance and reduce
    container footprint via multi-stage Alpine builds. The analysis below reflects their previous Python
    states.
5
6 ## Important Agent Instruction
7 - **NEVER use 'cat << EOF' or terminal redirections to write or modify files.** You must always use your
    native file-editing tools ('replace_file_content' or 'multi_replace_file_content'). Using 'cat EOF'
    leads to duplicated text, broken markdown headers, and structural corruption (which previously
    destroyed the section numbering in 'devref.md').
8 - **Always run 'bash scripts/crypto.sh --sign' after making any modifications to the core bash scripts.**
    This is required to update the cryptographic HMAC-SHA256 signatures; otherwise, the startup
    integrity checks will fail and block execution.
9 - **Always run 'git diff' and print the changes to the user *before* requesting or executing any git
    commit command.** The user must be shown the diff so they can review and understand the changes.
    Based on this diff, formulate a highly precise, detailed, and descriptive commit message (explaining
    exactly what was changed and why) rather than a generic summary, and present it to the user
    alongside the diff.
10 - **Always run 'git status' before 'git diff' to identify all modified and untracked files, ensuring no
    files (such as new scripts or configuration assets) are missed before formulating commit diffs and
    staging.**
11 - **Always use the '--restart' flag when asked to restart the toggle (e.g., run './toggle.sh --restart')
    .**
12 - **NEVER execute any 'sudo' commands (especially 'sudo route', 'sudo dscacheutil', or any network-
    modifying commands) without explicitly explaining to the user what the command does and why it is
    needed FIRST. Do not assume implicit permission. Wait for the user's approval before running it.**
13 - **When fixing an error and using logs to debug, always show the user the reasoning and the specific
    logs that support this theory.**
14 - **When the user tells you to "push", this means read git diffs and prepare commit messages that
    correspond to these changes (do not ignore any changes). If the changes can be atomized into
    multiple commits for easier understanding, do so.**
15 - **When the user says "latex this", it means you should run 'python3 scripts/generate_tex.py' to
    regenerate the unified LaTeX document.**
16 - **NEVER apply changes to running gateway containers directly via 'docker compose up' or 'docker compose
    restart'.** Recreating or restarting containers (especially the 'warp' container which hosts the
    network namespace) breaks the shared network stack for all other services ('adguardhome', 'tailscale',
    'routing-fix'), severing host connectivity and freezing macOS DNS. If you need to apply changes
    or rebuild containers, cleanly stop the gateway first using './toggle.sh', or trigger 'recover.sh --
    post-wake' to rebuild and re-verify the network stack safely in the correct dependency order. Do NOT
    run './toggle.sh --restart' or any network-rebuilding commands yourself if the execution could
    sever host network access; instead, explain the changes and ask the USER to run './toggle.sh --
    restart' (or './toggle.sh' stop and start) themselves to safely apply the modifications.
17
18 ## How to Verify Gateway is Working
19 Whenever modifications are made to the gateway scripts, routing, or containers, execute these
    verification checks in order:
20 1. **Container Health:** Run 'docker compose ps' to ensure all containers ('warp', 'adguardhome', '
    routing-fix', 'tailscale') are 'running' and 'healthy'.
21 2. **DNS Hijacking Status:** Run 'networksetup -getdnsservers "Wi-Fi"' (or appropriate network service)
    and ensure it is set exclusively to the gateway's Tailscale static IP (typically '100.100.21.8').
22 3. **DNS Query Interception:** Run 'dig @100.100.21.8 google.com' (using the gateway static IP from '
    gateway_ip') to ensure AdGuard Home is active, intercepting, and resolving queries.
23 4. **Double-Tunnel Egress:** Run 'curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp' to verify
    the egress is routed through Cloudflare WARP ('warp=on').
24 5. **Host Leak Monitoring:** Run 'tail -n 30 output.log' and verify the background watchers (DNS + WARP)
    report healthy status without 'LEAK' warnings. For a deeper on-demand audit, run 'bash scripts/
    diagnose-host-leak.sh' (add '--watch' to poll, '--fix' to remediate).
25
26 **Automated:** 'bash scripts/sweep.sh' runs a read-only 7-check health sweep covering the above plus the
    PF kill-switch, tailnet-wide peer reachability, and real throughput, writes a timestamped 'sweep-<
    UTC>.txt' report to the repo root, and exits non-zero on FAIL ('--quiet' for just the verdict). It
    never mutates routing/PF/DNS/containers, so it is safe against a live gateway. See its entry in Part
    3.

```



```

27 ---
28 ---
29
30 ## Part 1: macOS Main Scripts
31
32 ### toggle.sh (2191 lines, unified macOS + Linux)
33
34 > **Refreshed 2026-07-15.** Post-unification ('$OSTYPE' switch at L20) + the July-2026 'common.sh'
    extraction, the tables below reflect the *current* file. Most former "toggle.sh functions" now live
    in 'scripts/common.sh'; the host-leak-probe pair was removed. For the ordered execution flow see **
    toggle.sh Execution Map & Critical Path*** above.
35
36 **Functions still defined in 'toggle.sh'** several exist in *both* halves (Linux L63848 / macOS L8782191
    ):
37
38 | Function | Linux L | macOS L | Purpose |
39 |-----|-----|-----|-----|
40 | '_log_exit_breadcrumb()' | 30 | 885 | Write lifecycle-phase breadcrumb on exit (post-mortem trail) |
41 | 'get_free_port()' | 67 | 944 | Pick a random unused port (self- + kernel-collision checked) |
42 | 'run_gui_cmd()' | 1033 | | Run a command as the logged-in console GUI user |
43 | 'write_gateway_active_marker()' | 1053 | | Write '/tmp/nullexit-gateway-active.marker' (macOS only) |
44 | 'clear_gateway_active_marker()' | 1059 | | Remove active marker + 'TUNNEL_FAILED_CLOSED.marker' |
45 | 'start_sleep_prevention()' | 145 | 1091 | Start 'caffeinate' w/ revert trap |
46 | 'stop_sleep_prevention()' | 188 | | Stop caffeinate via PID file |
47 | 'stop_dns_watcher()' | 202 | | Stop DNS watcher (macOS uses 'stop_pidfile_daemon') |
48 | 'start_warp_watcher()' | 1164 | | Background WARP liveness monitor 'recover.sh' after N fails |
49 | 'stop_warp_watcher()' | 1241 | | Stop WARP watcher via PID file |
50 | 'cleanup_handler()' | 215 | 1258 | Fail-closed teardown trap (guarded by 'SUCCESS_RUN') |
51 | 'restart_tailscaled_daemon()' | 327 | *(common)* | Restart tailscaled (Linux-half local copy; macOS
    uses common.sh) |
52 | 'disconnect_tailscale_host()' | 337 | *(common)* | 'tailscale down' w/ retry (Linux-half local copy;
    macOS uses common.sh) |
53 | 'setup_exit_node_routing()' | 1430 | | Override default route through the Tailscale utun interface |
54 | 'cleanup_network_state()' | 354 | 1464 | Nuclear cleanup: proxies, DNS cache, routes, Wi-Fi, sharing
    services |
55
56 **Delegated to 'scripts/common.sh'** (moved out of toggle.sh in the July-2026 refactor; line = common.sh)
    :
57
58 | Function | common.sh L | Function | common.sh L |
59 |-----|-----|-----|-----|
60 | 'run_with_timeout()' | 58 | 'enable_socks_proxy()' | 769 |
61 | 'is_gateway_active()' | 168 | 'disable_socks_proxy()' | 803 |
62 | 'gateway_state()' | 232 | 'reset_dns()' | 824 |
63 | 'get_active_service()' | 284 | 'stop_local_dns_proxy()' | 898 |
64 | 'restart_tailscaled_daemon()' | 338 | 'start_local_dns_proxy()' | 909 |
65 | 'disconnect_tailscale_host()' | 352 | 'force_dns_to_gateway()' | 988 |
66 | 'stop_pidfile_daemon()' | 393 | 'add_warp_bypass_routes()' | 484 |
67 | 'start_dns_watcher()' | 715 | 'remove_warp_bypass_routes()' | 551 |
68
69 **Removed:** 'start_host_leak_probe()' / 'stop_host_leak_probe()' and the whole 'HOST_LEAK_PROBE' prober
    subsystem no longer defined or called anywhere in 'toggle.sh', 'recover.sh', or 'common.sh', and '
    HOST_LEAK_PROBE' is no longer an env var (all verified). The only surviving host-leak tool is the
    standalone, on-demand 'scripts/diagnose-host-leak.sh' (there is **no** 'scripts/host-leak-probe.sh')
    .
70
71 **Constants / anchors (current):**
72
73 | Constant | Value | Line |
74 |-----|-----|-----|
75 | 'LOG_FILE' | '$PWD/output.log' | L22 (Linux) / L863 (macOS) |
76 | 'LOCK_FILE' | '/tmp/nullexit-toggle.lock' | L1020 |
77 | 'PID_FILE' (caffeinate) | '/tmp/nullexit-caffeinate.pid' | L142 |
78 | 'DNS_WATCHER_PID_FILE' | '/tmp/nullexit-dns-watcher.pid' | L200 |
79 | 'WARP_WATCHER_PID_FILE' | '/tmp/nullexit-warp-watcher.pid' | L1144 |
80 | Gateway active marker | '/tmp/nullexit-gateway-active.marker' | L1054 (write) |
81 | TUNNEL_FAILED_CLOSED marker | '<repo>/TUNNEL_FAILED_CLOSED.marker' | L1061 |
82 | 'SOCKS_PROXY_PORT' | random per-toggle, '1080' fallback | L91/L107 (Linux) L968/L986 (macOS) |
83 | DNS search domain | 'ts.net' | L763 |
84 | 'PATH' (via 'setup_standard_path', common.sh) | Homebrew + system bins | L1330 |
85 | WARP endpoints (via 'get_warp_endpoint_1/2', common.sh) | '162.159.192.1' / '162.159.193.1' | common.sh
    L335-336 |
86 | DNS fallback | '1.1.1.1' (in 'reset_dns', common.sh) | common.sh L824 |
87
88 ---
89
90 ### toggle.sh Execution Map & Critical Path (verified 2026-07-15)
91
92 > Traced from the current unified 'toggle.sh' (2191 lines). The file carries **two halves** selected by '
    $OSTYPE' at **L20** ('if [[ "$OSTYPE" != "darwin"* ]] runs the self-contained **Linux path L63848
    ** then 'exit 0' at ~L845; darwin falls through to the **macOS path L8782191**). Line numbers below

```

```

    are the **macOS path** (the deployment target). The halves are *similar in intent but not identical
    * see "Linux path where it diverges" below. '~Line' = current, approximate.
93
94 ##### Entry & dispatch
95
96 '''
97 ./toggle.sh [ (no-arg) | --restart | restart ]
98 |
99 | -- verify HMAC signatures (crypto.sh --verify) --fail--> exit 1      (L857)
100 | -- rotate output.log if > 1GB
101 | -- arm teardown trap: ERR/INT/TERM/HUP -> cleanup_handler          (L1317-1320)
102 | -- generate random ports (Tor / SOCKS / DNS)                        (L940)
103 |
104 | -- --restart --> gateway_state? running -> run STOP, then START    (L999)
105 |
106 | -- no-arg --> gateway_state?                                       (L1535)
107 |     running -> STOP branch (L1538-1584)
108 |     stopped -> START branch (L1590-2191)
109 '''
110
111 'gateway_state' is driven by '/tmp/nullexit-gateway-active.marker' (persisted intent), falling back to a
    live container probe.
112
113 ##### START synchronous critical path (all must complete)
114
115 | # | Step | ~Line | Effect if it fails / is interrupted |
116 | ---|-----|-----|-----|
117 | 1 | 'disconnect_tailscale_host' (clean slate) | 1595 | host left mesh |
118 | 2 | 'colima_start_until_ready' (readiness poll, 15.8.7) | 1631 | no Docker VM abort |
119 | 3 | honey-port reap ('pkill') + arm ('nc -l', disowned) | 1695 | tripwire missing (non-fatal) |
120 | 4 | rule-compile (hash-gated, OFF critical path 15.8.8) | 1708 | reuse prior compiled rules |
121 | 5 | build-gate routing-fix / tor images | 1743 | rebuild only on change |
122 | 6 | 'docker compose up -d' warp, tailscale, routing-fix, adguard, tor | 1758 | no gateway |
123 | 7 | wait tailscaled reachable | 1907 | host tailscale unusable |
124 | 8 | **host 'tailscale set --exit-node=$TS_IP'** | 2046 | **host not routed via gateway REAL-IP LEAK**
125 | 9 | 'add_warp_bypass_routes' + 'setup_exit_node_routing' | 2048 | WARP endpoint unreachable / wrong
    route |
126 | 10 | **force_dns_to_gateway** (+ 'start_local_dns_proxy') | 2078 | host DNS not hijacked |
127 | 11 | 'enable_socks_proxy' + verify SOCKSWARP | 2100 | SOCKS off |
128 | 12 | print **"STARTED in Ns" (timing marker) | 2127 | cosmetic only NOT the commit point |
129 | 13 | 'start_sleep_prevention' (caffeinate) | 2130 | Mac may sleep |
130 | 14 | 'start_dns_watcher' (re-hijack DNS every 30s) | 2134 | no roam DNS recovery |
131 | 15 | 'start_warp_watcher' (warp liveness recover.sh) | 2138 | no WARP self-heal |
132 | 16 | **write_gateway_active_marker** | 2146 | watcher believes gateway is DOWN |
133 | 17 | **SUCCESS_RUN=true** disarms the teardown trap | 2164 | *(true commit point)* |
134 | 18 | print "You can close this terminal window now." | 2191 | |
135
136 ##### STOP critical path
137
138 'disconnect_tailscale_host' (1543) 'docker compose down --remove-orphans' (1552) 'colima stop' (1558)
    'reset_dns' (1568) 'stop_local_dns_proxy' (1571) stop caffeinate (1574) stop DNS watcher (1575)
    'stop_warp_watcher' (1576) 'clear_gateway_active_marker' (1577) "STOPPED in Ns" (1584).
139
140 ##### Long-lived background processes spawned by START
141
142 | Process | Started | PID / reap | Role |
143 | ---|-----|-----|-----|
144 | honey-port 'nc -l localhost $PORT' (disowned) | L1695 | 'pkill -f "nc -l localhost"' | port-scan
    tripwire |
145 | 'caffeinate' | L2130 | '/tmp/nullexit-caffeinate.pid' | prevent system sleep |
146 | DNS watcher loop | L2134 | '/tmp/nullexit-dns-watcher.pid' | re-hijack DNS every 30s (Wi-Fi roam) |
147 | WARP watcher loop | L2138 | '/tmp/nullexit-warp-watcher.pid' | monitor warp; fire recover.sh after N
    fails |
148 | local DNS proxy 'dns-proxy.py' | L2089 | (killed by 'stop_local_dns_proxy') | UDP:53 gateway DNS
    fallback |
149
150 Separately, the **launchd 'scripts/watcher.sh' daemon (post-wake / post-roam recovery) runs
    independently of toggle and invokes 'recover.sh --post-wake'; toggle does **not** spawn it.
151
152 ##### 'cleanup_handler' the fail-closed teardown trap
153
154 Armed on 'ERR/INT/TERM/HUP' (L1317-1320), guarded by 'SUCCESS_RUN'. While 'SUCCESS_RUN=false' it runs: '
    reset_dns', 'disconnect_tailscale_host', 'stop_local_dns_proxy', stop watchers, '
    clear_gateway_active_marker', 'docker compose down' (+ 'colima stop' on 'ERR'). This is the **fail-
    closed guarantee**: a half-built gateway is torn down, not left leaking.
155
156 ##### Correctness findings (verified against source)
157
158 1. **"STARTED in Ns" (L2127) is NOT the commit point 'SUCCESS_RUN=true' (L2164) is.** The marker write (
    L2146) and watcher starts (L2130-2138) sit *between* them. **Any TERM/INT including a wrapping tool

```



```

's timeout in the L2127L2164 window fires 'cleanup_handler' a FULL teardown** (containers + colima
+ host tailscale), leaving the marker ABSENT and the host on its raw ISP link, *despite* the
printed success. This is exactly the 2026-07-15 incident. (The Linux half has the same shape but a *
much* narrower window: STARTED L814 'SUCCESS_RUN=true' L820, and it maintains no marker.)
159 2. **Never pipe 'toggle.sh' through a reader** ('| tail', '| grep'). The disowned children (honey-port,
watchers, dns-proxy, caffeinate) inherit stdout, so the pipe never EOFs and the reader hangs long
after toggle has exited which then trips the wrapper's timeout SIGTERM finding #1. Run detached
with a file redirect instead: 'nohup ./toggle.sh --restart > /tmp/t.log 2>&1 &', then poll the file.
160 3. **Host wiring is correctly on the synchronous critical path** (steps 8 & 10 precede "STARTED"), so a *
completed* START cannot leave the host unrouted/leaking the only way to that state is an
interruption per finding #1.
161 4. **Namespacing ordering:** step 6's 'docker compose up' creates warp's network namespace that 'tailscale
'/'routing-fix'/'adguard' share (compose 'depends_on: warp healthy' enforces order). Never 'docker
compose restart warp' on a live gateway it detaches the shared netns (see Agent Instruction, L16).
162
163 ##### Linux path where it diverges (L63848)
164
165 The Linux half is intent-similar but **not** a line-for-line mirror of macOS. Verified differences:
166
167 | Aspect | macOS (L8782191) | Linux (L63848) |
168 |-----|-----|-----|
169 | VM layer | Colima ('colima_start_until_ready', L1631) | none native Docker |
170 | Host exit-node assert | '$TS_BIN set --exit-node="$TS_IP"' (L2046) | '$TS_BIN up --reset --exit-node
="$TS_IP"' (L735) |
171 | gateway-active marker | written (L2146) | **not maintained** (L119) 'gateway_state' degrades to a live
container probe |
172 | Commit window | STARTED L2127 'SUCCESS_RUN' L2164 | STARTED L814 'SUCCESS_RUN' L820 |
173 | Port gen tooling | 'jot' / 'netstat' | 'shuf'/'$RANDOM' / 'ss' |
174
175 The 'up --reset' vs 'set' distinction matters: '--reset' re-applies *all* prefs and forces a route re-
eval, whereas 'set' mutates only the exit-node pref the same asymmetry 'recover.sh --post-wake'
relies on (devref 15.12.21).
176
177 ---
178
179 ### recover.sh (566 lines, 23KB)
180
181 **Functions (7 total):**
182
183 | # | Function | Lines | Purpose |
184 |---|---|---|---|
185 | 1 | 'step()' | L40 | Print bold step header |
186 | 2 | 'ok()' | L41 | Print green success message |
187 | 3 | 'warn()' | L42 | Print yellow warning message |
188 | 4 | 'fail()' | L43 | Print red failure message |
189 | 5 | 'die()' | L44 | Print red error and exit |
190 | 6 | 'run_with_timeout()' | L5674 | Run command with timeout (bash-native, no GNU coreutils) |
191 | 7 | 'restart_tailscaled_daemon()' | L118129 | Restart tailscaled daemon via brew services |
192
193 **Constants:**
194
195 | Constant | Value | Line |
196 |-----|-----|-----|
197 | Color codes | 'RED', 'GREEN', 'YELLOW', 'BOLD', 'NC' | L37-38 |
198 | 'SCRIPT_DIR' | '${dirname "${BASH_SOURCE[0]}"}' | L50 |
199 | 'PATH' export | '/opt/homebrew/bin:/usr/local/bin:$PATH' | L77 |
200 | 'PID_FILE' | '/tmp/nullexit-caffeinate.pid' | L387 |
201 | 'DNS_WATCHER_PID_FILE' | '/tmp/nullexit-dns-watcher.pid' | L407 |
202 | WARP endpoints | '162.159.192.1', '162.159.193.1' | L329-336 |
203 | TUNNEL_FAILED_CLOSED marker | '$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker' | L51 |
204
205 **Duplicated logic from toggle.sh:**
206 *(Note: As of July 2026, **ALL** of the below duplicated logic has been successfully extracted to '
scripts/common.sh!)*
207 - 'run_with_timeout()' simpler subshell version vs toggle's PID-tracking version
208 - 'restart_tailscaled_daemon()' near-identical (uses warn()/ok() vs echo)
209 - Active network service detection inline at L217-233 (toggle has 'get_active_service()' function)
210 - EN0 service resolution inline at L230-233 (same as toggle L515-516)
211 - WARP bypass routes inline at L329-337 (toggle has 'add_warp_bypass_routes()')
212 - Exit node routing setup inline at L340-345 (toggle has 'setup_exit_node_routing()')
213 - DNS cache flushing L296-302 (same as toggle L611-616)
214 - Wi-Fi power-cycling L442-461 (similar to toggle L626-637)
215 - Proxy disabling L282-288 (same as toggle L604-608)
216 - Sharing services reset L586 (same as toggle L642)
217 - PID file stop pattern L387-399, L407-419, L423-435 (same as toggle L157-167, L193-203, L312-322)
218 - KILL_SWITCH check L194 (same as toggle L141, L469)
219 - .gateway_ip reading L138-142, L209-213 (same pattern as toggle L1080)
220
221 ---
222
223 ### setup.sh (410 lines, 18KB)

```

```

224
225 **Functions (4 total):**
226
227 | # | Function | Lines | Purpose |
228 |---|-----|-----|-----|
229 | 1 | 'step()' | L10 | Print bold blue step header |
230 | 2 | 'ok()' | L11 | Print green success message |
231 | 3 | 'warn()' | L12 | Print yellow warning message |
232 | 4 | 'die()' | L13 | Print red error and exit |
233
234 **Constants:**
235
236 | Constant | Value | Line |
237 |-----|-----|-----|
238 | Color codes | 'RED', 'GREEN', 'YELLOW', 'BLUE', 'BOLD', 'NC' | L7-8 |
239 | 'RECOMMENDED_TS_VERSION' | '1.98.5' | L71 |
240 | 'ADGUARD_PASSWORD' | 'nullexit' | L215 |
241 | Default .env values | 'GATEWAY_RULE_PROFILE=medium', 'GATEWAY_MSS=1120', etc. | L240-248 |
242
243 ---
244
245 ## Part 2: Linux Scripts
246
247 ### toggle.sh Linux path (unified into the root cross-platform script)
248
249 The Linux toggle is no longer a separate file. 'scripts/toggle-linux.sh' was
250 merged into the repo-root 'toggle.sh', which dispatches on '$OSTYPE': an early
251 'if [[ "$OSTYPE" != darwin* ]]' block runs the self-contained Linux toggle
252 (shuf / ss / nmcli / ip / systemd) and exits before the macOS body. The shared
253 DNS/proxy/daemon primitives it used were already moved to 'common.sh' in the
254 '313dc32' unification, so both branches call them by the same names.
255
256 ---
257
258 ### recover.sh Linux path (unified into the root cross-platform script)
259
260 Linux recovery is no longer a separate file. 'scripts/recover-linux.sh' was
261 merged into the repo-root 'recover.sh', which dispatches on '$OSTYPE': the Linux
262 branch runs a self-contained teardown (resolvectl / nmcli / ip) and exits before
263 the macOS dual-mode body. All formatting/timeout helpers come from 'common.sh'.
264
265 ---
266
267 ### scripts/setup-linux.sh (416 lines, 18KB)
268
269 **Functions (4 total):**
270
271 | # | Function | Line | Purpose |
272 |---|-----|-----|-----|
273 | 1 | 'step()' | L10 | Print formatted step header |
274 | 2 | 'ok()' | L11 | Print green success message |
275 | 3 | 'warn()' | L12 | Print yellow warning message |
276 | 4 | 'die()' | L13 | Print red error + exit 1 |
277
278 **~95% identical to macOS setup.sh.** Only differences:
279 1. Desktop shortcuts (L366-391 Linux '.desktop' files vs macOS '.app' via 'osacompile')
280 2. Final instructions text
281 3. .env template: macOS adds 'GATEWAY_USE_EXIT_NODE', 'WARP_FAIL_THRESHOLD', 'KILL_SWITCH' Linux omits
   these
282
283 ---
284
285 ## Part 3: Utility Scripts
286
287 ### scripts/diagnose-host-leak.sh (722 lines, 38KB)
288
289 **Purpose:** Comprehensive host-routing diagnostic. 8 checks classify host IP leaks. Supports '--fix' and
   '--watch' modes.
290
291 **Duplicated logic:**
292 - 'resolve_active_service()' identical pattern as toggle.sh, recover.sh, fix-docker-bridge-collision.sh
293 - 'run_with_timeout()' reimplemented timeout
294 - Color codes RED/GREEN/YELLOW/BOLD/NC with tty detection
295 - WARP check via 'cloudflare.com/cdn-cgi/trace'
296 - 'export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"'
297 - .gateway_ip reading pattern
298
299 ### scripts/fix-docker-bridge-collision.sh (402 lines, 18KB)
300
301 **Purpose:** Fixes Docker bridge subnet collisions with '10.200.1.0/24'.
302

```

```

303 **Duplicated logic:**
304 - 'resolve_active_iface()' same interface detection as diagnose-host-leak.sh
305 - Color codes, step()/ok()/warn()/fail()/die() formatting helpers
306
307 ### scripts/device-scramble.sh (159 lines, 7.8KB)
308
309 **Purpose:** Randomizes hostname and MAC address for Wi-Fi anonymity. Includes fail-safe MAC test and
    disassociate-first logic.
310
311 **Duplicated logic:**
312 - 'need_root()' same sudo check as other scripts
313 - 'WARP_INHIBIT' marker interacts with watcher.sh to prevent gateway shutdown during Wi-Fi blip
314
315 ### scripts/fix-ssh-delay.sh (63 lines, 2.7KB)
316
317 **Purpose:** One-shot fix for ~20s SSH delay. Standalone, minimal duplication (uses hardcoded / instead
    of color functions).
318
319 ### scripts/reinstall-host-tailscale.sh (740 lines, 31KB)
320
321 **Purpose:** Complete uninstall + reinstall of host Tailscale.
322
323 **Duplicated logic:**
324 - 'with_timeout()' yet another bash timeout implementation (polling-based, different from diagnose and
    toggle versions)
325 - Color codes + logging functions
326 - System Extension detection logic (similar to diagnose-host-leak.sh)
327
328 ### scripts/routing-fix.sh (266 lines, 10.5KB)
329
330 **Purpose:** Runs INSIDE warp container. Maintains routing table 200, FORWARD rules, country blocking, IP
    blocklists, tunnel health checks.
331
332 **Duplicated logic:**
333 - (Fixed) WARP_ENDPOINT defaults are now centralized in common.sh
334 - WARP health check via cdn-cgi/trace (another instance)
335 - iptables dual-backend pattern (for-loop over 'iptables iptables-legacy')
336
337 ### scripts/sweep.sh (491 lines, 28KB)
338
339 **Purpose:** Read-only post-restart health sweep answers "after a toggle/restart/wake, is the gateway
    actually healthy end-to-end including under real load?" Never mutates routing, PF, DNS, or
    containers, so it is safe to run against a live gateway at any time (including from a remote session
    ). 'errexit' is deliberately OFF: probes are *expected* to fail on an unhealthy gateway and each
    exit status is recorded as the diagnostic signal.
340
341 **The 7 checks** (each recorded PASS/WARN/FAIL; exit 0 only when nothing FAILs, so it can gate CI/launched
    /post-restart hooks):
342 1. **Containers** 'docker compose ps': warp/adguardhome/tailscale/tor/routing-fix up (warp + adguardhome
    additionally 'healthy').
343 2. **Double-tunnel/leak** container WARP egress IP == host egress IP, both 'warp=on' (gluetun ships wget
    not curl, so the in-container trace falls back to wget).
344 3. **Hostcontainer path** the gateway peer line in 'tailscale status' shows 'direct <ip:port>'; WARNs on
    a DERP relay.
345 4. **AdGuard filtering** 'compiled_rules.txt' block counts present + live interception ('dig @' the
    gateway IP from '.gateway_ip').
346 5. **PF kill-switch** 'pfctl' reports Enabled AND anchor 'com.apple/nullexit' carries rules (needs
    passwordless sudo, else WARN).
347 6. **Tailnet reachability** (added 2026-07-17) every peer 'tailscale status' reports online (self and '
    offline' peers skipped; a '-' status means online-but-idle and is kept) must answer 'tailscale ping
    --timeout=3s --c 1 --until-direct=false' (TSMP probes the WireGuard data path even on devices that
    drop ICMP). All pong PASS; some unreachable WARN; none FAIL. DERP-relayed pongs count as
    reachable (direct-only enforcement stays in check 3, which covers the gateway); DERP-relayed peers
    are also flagged inline as throughput-degraded, pointing at check 7.
348 7. **Throughput** (added 2026-07-18) pulls a real file over HTTPS from a fast CDN (Hetzner; Cloudflare's
    own speed-test endpoint 403s WARP's shared/CGNAT egress IPs, so it's avoided) twice: once over the
    host's normal double-tunnel route, once through the 'warp' container's own SOCKS5 proxy (port read
    from the live '/tmp/nullexit-ports.env') to isolate WARP from the Tailscale-wrap hop without exec'
    ing into or installing anything in the container. Verdict is read from curl's actual exit code, not
    a byte-count guess: '56' ('Recv failure: Connection reset by peer') is a real mid-transfer stall (
    WARN see devref.md 15.3.1); '28' is just the bounded test window ending on an otherwise-healthy
    transfer (PASS). A slow but supposedly "fast" CDN target was deliberately avoided after 'speedtest.
    tele2.net' (the initial choice) turned out to be the bottleneck itself, not nullexit see devref.md
    15.3.6.
349
350 **Output:** coloured stdout summary + a timestamped plain-text report 'sweep-<TIMESTAMP>.txt' in the repo
    root (one file per run, so runs can be diffed); '--quiet' prints only the verdict. Signed in '.
    signatures' re-run 'bash scripts/crypto.sh --sign' after editing it.
351
352 ### scripts/watcher.sh (269 lines, 13KB)
353

```

```

354 **Purpose:** Long-running launchd daemon for post-wake/post-roam recovery.
355
356 **Key function 'detect_lan_p2p_mode()':**
357 Added July 10, 2026. Called on startup and on every network state change (Listener 2 NET: events).
    Determines whether the current Wi-Fi network safely allows direct Tailscale LAN P2P connections:
358 1. **WPA2-Enterprise check:** 'system_profiler SPAirPortDataType' reads the 'Security:' field for the
    current association. Enterprise = 802.1x = always AP-isolated sets 'allow=false'.
359 2. **AP isolation probe:** 'route -n get 1.1.1.1' finds default gateway IP, then 'ping -c 3 -t 1 -q "$gw
    "' if 0 responses on a non-empty network, AP isolation is active sets 'allow=false'.
360 3. **Output:** Writes 'true' or 'false' to '.lan_p2p_detected' in the repo root.
361 4. **Consumer:** 'routing-fix.sh' reads this file every 30 seconds and enforces/relaxes the RFC1918 DROP
    rule accordingly no restart required.
362
363 **Duplicated logic:**
364 - PATH export (similar to toggle.sh/recover.sh)
365 - Marker file pattern ('/tmp/nullexit-gateway-active.marker')
366
367 ### scripts/dns-proxy.py (92 lines, 2.9KB) Clean, standalone
368
369 ### scripts/logger.py (144 lines, 5.9KB) Clean, standalone
370
371 ---
372
373 ## Part 4: Cross-Cutting Issues
374
375 ### Functions Identical Across macOS Linux
376
377 | Function | macOS toggle | macOS recover | macOS setup | Linux toggle | Linux recover | Linux setup |
378 |-----|-----|-----|-----|-----|-----|-----|
379 | 'run_with_timeout()' | | | identical | identical | | |
380 | 'stop_local_dns_proxy()' | | | identical | | | |
381 | 'start_local_dns_proxy()' | | | ~95% similar | | | |
382 | 'is_gateway_active()' | | | ~80% similar | | | |
383 | 'step/ok/warn/die' | | | identical | identical | | |
384
385 ### Platform-Adapted Functions (same logic, different OS commands)
386
387 | Function | Linux Command | macOS Command |
388 |-----|-----|-----|
389 | 'get_active_service()' | 'ip route get 1.1.1.1 \\\ awk '{print $5}' | 'route get default' + '
    networksetup -listnetworkserviceorder' |
390 | 'reset_dns()' | 'resolvectl revert' | 'networksetup -setdnsservers' |
391 | 'cleanup_network_state()' | 'ip link set dev down/up' | 'networksetup -setairportpower off/on' + proxy
    disable |
392 | 'force_dns_to_gateway()' | 'resolvectl dns' + 'resolvectl domain ~ts.net' | 'networksetup -
    setdnsservers' + '-setsearchdomains ts.net' |
393 | 'start_sleep_prevention()' | 'systemd-inhibit --what=sleep' | 'caffeinate -i' |
394 | 'enable/disable_socks_proxy()' | Stub (no-op on Linux) | 'networksetup -setsocksfirewallproxy' |
395 | DNS cache flush | 'resolvectl flush-caches' | 'dscacheutil -flushcache' + 'killall -HUP mDNSResponder'
    |
396
397 ### Features Missing From Linux Scripts
398
399 These exist in 'toggle.sh's macOS branch but NOT its Linux branch:
400
401 1. 'write_gateway_active_marker()' / 'clear_gateway_active_marker()'
402 2. 'restart_tailscaled_daemon()'
403 3. 'start_warp_watcher()' / 'stop_warp_watcher()' WARP liveness monitor
404 5. 'add_warp_bypass_routes()' / 'remove_warp_bypass_routes()' / 'setup_exit_node_routing()'
405 6. 'LOG_FILE' structured logging (timestamps to output.log)
406 7. MSS clamping injection (step 4c)
407 8. Rule count reporting
408 9. '--post-wake' mode in recover
409 10. Container teardown on failure in cleanup_handler
410
411 ---
412
413
414
415 ## Part 5: Constant Duplication Map
416
417 | Value | Files |
418 |-----|-----|
419 | '162.159.192.1' (WARP endpoint) | docker-compose.yml, routing-fix.sh (now dynamically loaded via .env/
    common.sh!) |
420 | '5335' (AdGuard DNS port) | docker-compose.yml, AdGuardHome.yaml, post-rules.txt |
421 | '5354' (mapped DNS port) | docker-compose.yml, dns-proxy.py |
422 | '41642' (Tailscale WireGuard) | docker-compose.yml, routing-fix.sh, post-rules.txt |
423 | '1080' (SOCKS5 port) | docker-compose.yml, toggle.sh |
424 | '1120' (MSS clamp) | post-rules.txt, .env |
425 | 'admin:nullexit' (AdGuard creds) | AdGuardHome.yaml (bcrypt) |

```

```

426 | '10.200.1.0/24' (Docker subnet) | (Fixed) no longer duplicated |
427 | 'America/New_York' (timezone) | docker-compose.yml (4 services) |
428 | '/tmp/nullexit-caffeinate.pid' | (Fixed) centralized in common.sh |
429 | '/tmp/nullexit-dns-watcher.pid' | (Fixed) centralized in common.sh |
430 | '/opt/homebrew/bin:... ' (PATH) | (Fixed) centralized in common.sh |
431
432 ---
433
434 ## Part 6: Refactoring Outcome (Implemented July 2026)
435
436 **Decision: "Scripts-only" architecture**
437 Instead of introducing a new 'lib/' directory, all shared code was moved into the existing 'scripts/'
    directory to keep the project hierarchy flat and simple.
438
439 ### 1. 'scripts/common.sh' (288 lines, 10KB)
440 Extracts the fundamental primitives and networking logic used across orchestrator scripts:
441 - **Formatters:** 'step()', 'ok()', 'warn()', 'fail()', 'die()'
442 - **Timeout Wrapper:** 'run_with_timeout()' (Canonical version with PID tracking, graceful SIGTERM, and
    SIGKILL fallback)
443 - **Daemon Lifecycle:** 'restart_tailscaled_daemon()', 'stop_pidfile_daemon()' (collapsed caffeinate, DNS
    watcher, and WARP watcher teardown logic)
444 - **Network Interface/Routing:** 'get_active_service()', 'get_en0_service()', 'bounce_wifi_interfaces()',
    'add_warp_bypass_routes()', 'remove_warp_bypass_routes()'
445 - **Proxy/DNS Cleanup:** 'disable_all_proxies()', 'flush_dns_cache()', 'reset_sharing_services()'
446 - **Configuration Helpers:** 'read_adguard_ip()', 'is_kill_switch_enabled()', 'get_warp_endpoint_1()', '
    get_warp_endpoint_2()' (which centralizes the hardcoded 162.159.192.1 / .193.1 into .env logic)
447
448 ### 2. 'scripts/setup-common.sh' (~350 lines)
449 Extracts the heavy, duplicated installation logic shared between 'setup.sh' (macOS) and 'scripts/setup-
    linux.sh':
450 - Docker and Colima prerequisite checks
451 - 'wgcf' binary download and WARP key generation
452 - '.env' configuration file generation
453 - Interactive Tailscale Auth Key and AdGuard Rule Profile prompts
454
455 ### Sourcing Pattern
456 - macOS root scripts ('toggle.sh', 'recover.sh', 'setup.sh') source via:
457   'source "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/common.sh"'
458 - Scripts under 'scripts/' ('watcher.sh', 'diagnose-host-leak.sh', etc.) source via:
459   'source "$(dirname "${BASH_SOURCE[0]}")/common.sh"'
460
461 ## Agent Verbs / Protocols
462
463 ### '/sweep' (or 'sweep the gateway')
464 When the user invokes this verb, the agent MUST perform a full end-to-end connectivity, health, security,
    and performance audit of the gateway:
465 1. **WARP Validation:** Run 'curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp=' (must equal '
    on').
466 2. **Tor Proxy & Control Validation:**
467   - Source the dynamically generated ports from '/tmp/nullexit-ports.env' (or identify them using '
    docker compose ps').
468   - **SOCKS5 Check:** Run 'curl -s --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.
    org/api/ip' to confirm the SOCKS5 proxy successfully connects and routes through the Tor network.
469   - **Control Port Check:** Query the Tor Control Port using netcat: 'echo "PROTOCOLINFO" | nc
    127.0.0.1:$TOR_CONTROL_PORT'. Verify it returns a standard '250-PROTOCOLINFO' response, confirming
    the Control Port is active and communicating.
470   - **Circuit & Node Lookup:** Query the Tor Control Port for the active path ('GETINFO circuit-status')
    . For each active built circuit, parse the relay fingerprints, lookup their IP addresses ('GETINFO
    ns/id/<fingerprint>'), and resolve their country codes using Tor's local GeoIP database ('GETINFO ip
    -to-country/<IP>'). Include the list of active circuits, showing the role (Guard/Middle/Exit),
    nickname, IP, and country of each hop in the final sweep report.
471   - **Transparent Onion Validation:** Run 'nslookup
    duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion' to verify it resolves to an IP
    within '198.18.0.0/15' benchmark range. Verify transparent dark web routing by running 'curl -sI -H
    "Host: duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion" http://<resolved-ip>/' (
    should return HTTP 301/200).
472 3. **DNS Leak Validation:**
473   - Audit the upstream resolver currently utilized by the host. Run 'curl -s https://edns.ip-api.com/
    json' to fetch details of the DNS resolver processing the client's requests.
474   - Verify that the resolver IP and organization returned belong to Cloudflare (e.g., AS13335) or match
    the WARP tunnel's exit range, and that **no** DNS traffic leaks to the user's raw physical ISP DNS.
475 4. **Performance Benchmark Audit:**
476   - Run the python benchmark script 'python3 benchmarks/test_load.py'.
477   - Report the overall load speed and the ratio of successfully fetched vs. blocked/failed subresources
    for the targeted domains to ensure ad-blocking and MTU encapsulation are performing correctly.
478 5. **Tailscale Mesh Validation:**
479   - Run 'ping -c 1 100.100.21.8' to ensure the gateway is reachable.
480   - Run 'tailscale status' to identify all active/online nodes in the mesh.
481   - For every online peer returned in the status output, execute 'tailscale ping --timeout=3s --c 1 --
    until-direct=false <peer-ip>' (TSMP probes the WireGuard data path and works even on phones that
    drop ICMP) to verify connectivity and confirm healthy routing to all active mesh devices. 'bash

```

```

scripts/sweep.sh' automates this as its check 6/7 and reports each peer's path (direct endpoint vs '
DERP(...)') and latency; its check 7/7 additionally measures real throughput (host path vs. a WARP-
only baseline).
482 6. **Log Audit:** Run 'tail -n 100 output.log | grep -iE "(error|fail|warn)"' to catch any underlying
component failures or warnings.
483
484 ### 'latex' (or 'generate latex')
485 When the user invokes this verb, the agent MUST run 'python3 -I scripts/generate_tex.py' to regenerate
the documentation source code ('nullexit_unified.tex'). The agent MUST NOT attempt to compile the '.
tex' file into a PDF (the user handles PDF compilation locally).

```

4 devref.md

```

1 # nullexit Development Reference & Resolved Issues
2
3 > **Last updated:** July 12, 2026
4 > **Purpose:** Provide any LLM or developer with complete project understanding, debugging history, and
resolved issues so they can make informed changes without re-reading every file.
5 > **Diagrams:** See ['diagrams.md'](./diagrams.md) for system architecture, toggle.sh flowcharts,
monitoring layer, traffic sequence, recover.sh decision tree, and the full failureself-healing map.
6
7 This reference is organized into four parts:
8 > - **Part I Architecture & Reference** (18)
9 > - **Part II Design Decisions & Threat Model** (914)
10 > - **Part III Incident Log & Resolved Issues** (15)
11 > - **Part IV Observations, Changelog & TODO** (1618)
12
13 ---
14
15 # Part I Architecture & Reference
16
17 ## 1. What Is nullexit?
18
19 nullexit is a **Tunnel-in-Tunnel VPN gateway** that chains **Tailscale** (mesh VPN) through **Cloudflare
WARP** (exit VPN). It runs on a macOS host inside Docker containers managed by **Colima** (
lightweight Linux VM). The goal: any device on the user's Tailscale mesh can use this gateway as an
exit node, and all traffic exits through Cloudflare WARP achieving double encryption, ISP
invisibility, and network-wide ad-blocking via **AdGuard Home**.
20
21 ### Traffic Flow
22 ```
23 Phone/Laptop Tailscale mesh (WireGuard) Mac host Docker container
24 Tailscale decrypts Gluetun re-encrypts WARP tun0 Cloudflare Internet
25 ```
26
27 ### SOCKS5 Failover Path (if exit node is unavailable)
28 ```
29 DNS: Browser host 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP DNS
30 TCP: Apps macOS SOCKS5 proxy (127.0.0.1:1080) container tun0 WARP Internet
31 ```
32
33 ---
34
35 ## 2. Architecture
36
37 ### Containers (all share 'warp's network namespace via 'network_mode: service:warp')
38 | Container | Image | Role |
39 |-----|-----|-----|
40 | 'warp' | 'qmcgaw/gluetun:v3.41.1' | Gluetun WireGuard client Cloudflare WARP. Owns the network
namespace. Strict firewall. |
41 | 'tailscale' | 'tailscale/tailscale:v1.98.4' | Advertises as exit node on the mesh. Also provides the
built-in SOCKS5 proxy fallback via 'TS_SOCKS5_SERVER=:1080'. |
42 | 'routing-fix' | 'alpine:3.20' | Sidecar that maintains routing tables + iptables rules every 30 seconds
. |
43 | 'rule-compiler' | 'golang:1.22-alpine' / 'alpine:3.20' | One-shot startup container that compiles 500k+
DNS block rules in <2s and immediately exits. |
44 | 'adguardhome' | 'adguard/adguardhome:v0.107.77' | DNS sinkhole for ads/trackers. Listens on port 5335.
Upstream DNS: '127.0.0.1:53' (through WARP). |
45
46 ### Port Mappings (on host via Colima)
47 | Host Port | Container Port | Protocol | Purpose |
48 |-----|-----|-----|-----|
49 | 5354 | 5335 | TCP+UDP | AdGuard DNS |
50 | (Tailscale IP):3000 | 3000 | TCP | AdGuard web UI (Internal to Mesh) |
51 | 41641 | 41641 | UDP | Tailscale WireGuard direct |
52 | 1080 | 1080 | TCP | SOCKS5 proxy |

```



```

53
54 ### Host Components
55 - **Colima VM** Linux VM running Docker (600MB RAM + 400MB swap)
56 - **tailscaled** Standalone Tailscale daemon via 'brew install tailscale' + 'brew services start
    tailscale' (NOT the GUI '.app')
57 - **macOS network settings** DNS, SOCKS proxy, routing all manipulated by 'toggle.sh'
58
59 ---
60
61 ## 3. Key Files
62
63 | File | Lines | Purpose |
64 |-----|-----|-----|
65 | 'toggle.sh' | 2109 | **Main script (macOS + Linux).** Dispatches on '$OSTYPE' an early Linux block (
    shuf/ss/nmcli/ip, systemd) runs and exits before the macOS body (jot/netstat/networksetup, launchd/
    pf). Detects state, toggles gateway ON/OFF. Handles DNS hijacking, Tailscale exit node, SOCKS proxy,
    sleep prevention, timeouts, cleanup. |
66 | 'setup.sh' | 406 | **One-time setup.** Installs deps (Docker, Tailscale, wgcf), generates WARP keys,
    writes '.env', configures AdGuard via API, starts containers. |
67 | 'docker-compose.yml' | 194 | Service definitions for all 5 containers. |
68 | 'scripts/routing-fix.sh' | 201 | Maintains routing tables (table 200 for SOCKS5, table 52 for CGNAT)
    and FORWARD iptables rules in both nftables and legacy backends. Loads and enforces the IP blocklist
    via 'ipset'. Runs in a 30-second loop. |
69 | 'post-rules.txt' | 18 | Gluetun iptables rules loaded at container start. FORWARD accept for tailscale0
    , NAT MASQUERADE on tun0, DNS redirect to port 5335, IPv6 drop, TCP MSS clamping. |
70
71 | 'scripts/dns-proxy.py' | 91 | Local DNS proxy. UDP:53 TCP:5354 with proper 2-byte length prefix.
    Fallback when exit node is unavailable. |
72 | 'scripts/rule-compiler/' | 800 | Unified threat compiler (ported to Go from Python). **DNS pipeline:**
    fetches remote blocklists, deduplicates against AdGuard native filters, optimizes subdomains (~50%
    reduction), compiles to AdGuard syntax. **IP pipeline:** fetches threat-intel feeds (Spamhaus, Feodo
    , ET, CINS), strips private/Tailscale ranges, outputs atomic 'ipset' restore file. Built dynamically
    in Docker multi-stage build. |
73 | **'adguard/work/'** | | **Bind-mounted container data folder. Off-limits from outside Docker READ-
    ONLY-EXCEPT-VIA-'sync-rules'. See README 6.** |
74 | 'recover.sh' | 742 | Nuclear recovery script (**macOS + Linux** dispatches on '$OSTYPE'; the Linux
    branch runs a self-contained teardown via 'resolvectl'/'nmcli'/'ip' and exits, macOS falls through
    to the dual-mode body). Gain: use '--post-wake' (macOS) for non-destructive refresh after sleep/wake
    or Wi-Fi roam keeps the gateway live while still re-hijacking DNS + refreshing Tailscale exit-node
    + force-recreating warp if unhealthy. Resets DNS on ALL services, disables proxies, flushes routes,
    stops containers, kills sleep prevention, power-cycles Wi-Fi. Cryptographically verified. |
75 | 'scripts/diagnose-host-leak.sh' | 580+ | **One-shot host-routing diagnostic.** Runs 8 checks (Tailscale
    , SOCKS5, CDN-cgi/trace, IPv6 leak, default route, output.log hints, gateway IP, Tailscale System
    Extension conflict), classifies into ONE of three known leak scenarios (A=SOCKS5 fallback, B=IPv6
    leak, C=route freeze) or OK, writes a timestamped report to 'host-leak-diagnostic-<UTC>.txt', and
    prints a ready-to-run fix on stdout. Pass '--fix' to apply the matched remediation and re-verify
    egress. Pass '--watch' (or '--watch 30') to run a full baseline then continuously monitor warp/IPv6/
    default-route every N seconds, alerting on any state change. See 15.6.1. |
76 | 'scripts/host-leak-probe.sh' | | **REMOVED (2026-07-15, verified deleted/untracked).** Was a
    continuous sub-second host-egress prober auto-launched via 'HOST_LEAK_PROBE'. Its fail-closed role
    is now covered by the **PF kill-switch** ('enable_killswitch' in 'common.sh') + the in-container **
    WARP Watcher** (15.6.2); on-demand auditing is 'scripts/diagnose-host-leak.sh'. See 15.6.1 (retained
    as historical rationale). |
77 | 'scripts/fix-docker-bridge-collision.sh' | 290 | **One-shot fix for 15.2.3 (Docker bridge subnet
    collision).** Detects Docker's '172.17.0.0/16' bridge colliding with the host Wi-Fi subnet (the
    symptom that produces 'default 172.17.0.1 UGScg en0' in 'netstat -rn'), picks a non-colliding '
    docker.bip' from a small candidate set, idempotently edits '~/.colima/default/colima.yaml' with a
    timestamped backup, restarts Colima, rebinds Wi-Fi DHCP, verifies the new default route is no longer
    Docker's bridge, and re-runs 'toggle.sh' + 'scripts/diagnose-host-leak.sh'. Supports '--dry' and
    '--skip-toggle'. |
78 | 'scripts/watcher.sh' | 194 | **Long-running post-wake + post-roam daemon.** Launched by launchd
    LaunchAgent; subscribes to 'com.apple.powermanagement' events via 'log stream' and to network-state
    changes via 'scutil n.watch'. Both fire 'recover.sh --post-wake'. Single-instance lock + SCRIPT_DIR-
    relative resolve of 'RECOVER'. See 15.5.1. |
79 | 'scripts/common.sh' | 40 | **Shared script primitives.** Centralizes formatted echo statements ('step
    ', 'ok', 'warn', 'die') and the safe background 'run_with_timeout' execution wrapper. Sourced by all
    orchestrator scripts across macOS and Linux. |
80 | 'scripts/setup-common.sh' | 350 | **Shared installation logic.** Centralizes logic for verifying
    Docker, installing Tailscale/wgcf, generating '.env', and prompting for auth keys. Sourced by 'setup
    .sh' and 'setup-linux.sh'. |
81 | 'scripts/setup-linux.sh' | 80 | **Linux sibling of 'setup.sh'.** Sources 'setup-common.sh'. Distro
    detection (apt/dnf/pacman), installs Linux-specific dependencies, creates '.desktop' shortcuts. |
82 | 'scripts/Toggle-Gateway.bat' | 300 | **Windows Toggle helper.** Moved from root to 'scripts/'. Helps
    invoke the framework on Windows if applicable. |
83 | 'scripts/reinstall-host-tailscale.sh' | 740 | **Nuke-and-reinstall tool for host Tailscale.**
    Completely purges Tailscale, its settings, macOS Background Items ('.btm' database), and stale
    System Extensions. Wipes ALL Login Items via 'sftool reset-login-items' (requires sudo). Useful for
    crisis recovery but destructive to other login items. |
84 | 'scripts/fix-ssh-delay.sh' | 100 | **Fixes SSH connection delays.** One-shot script to address SSH
    delays caused by DNS or routing issues, useful in crisis situations. |

```

```

85 | 'scripts/logger/' | ~100 | **Internal logger piped from 'scripts/routing-fix.sh'** inside the 'warp'
    container. Ported to Go from Python. Built via Docker multi-stage build. |
86 | 'black_list.txt' | 71 | Custom domains to block (ads, trackers, telemetry). Supports '$important'
    modifier. |
87 | 'white_list.txt' | 222 | Domains to force-allow (YouTube, Apple services, etc.). Always wins over
    blocks. |
88 | '.env' | ~14 | WARP WireGuard keys, Tailscale auth key, rule profile. **Contains secrets &
    NULLEXIT_SEED.** |
89 | '.gateway_ip' | 1 | Static gateway Tailscale IP (fallback for dynamic resolution). |
90 | 'scripts/unlock-files.sh' | ~20 | **One-shot stale-permission fix.** Uses atomic rename ('cp' + 'mv')
    to replace locked inodes (mode '000'/'0444' from old 'chmod' decisions) with fresh writable ones no
    'chmod' called. |
91 | 'scripts/crypt.sh' | ~50 | **Cryptographic integrity enforcement.** Uses 'NULLEXIT_SEED' from '.env'
    to sign core bash scripts (HMAC-SHA256) and strictly verify them on start to prevent manipulation.
    Pass '--sign' or '--verify'. |
92 | 'scripts/pf.conf' | 35 | **macOS native firewall ruleset.** Enforces the default-deny kill-switch,
    allows local LAN / Tailscale traffic, and applies TCP MSS Clamping ('max-mss 1160') to prevent
    WireGuard MTU fragmentation. |
93
94 ---
95
96 ## 4. toggle.sh Flow
97
98 ### State Detection
99 'is_gateway_active()' returns true if EITHER containers are running OR host DNS is not '1.1.1.1'. This
    dual check prevents the script from starting when it should stop (e.g., containers crashed but DNS
    is still hijacked).
100
101 ### START Path (gateway OFF ON)
102 1. **Reset DNS to 1.1.1.1** Prevents deadlocks during startup
103 2. **Disconnect host Tailscale** 'tailscale down' (prevents exit-node routing during container boot)
104 3. **Compile DNS rules** 'docker compose run --rm rule-compiler' (Go binary inside multi-stage Docker
    build)
105 4. **Boot Colima VM** (if not running) via 'colima_start_until_ready()', which returns as soon as Docker
    is reachable (~12s) instead of blocking on Colima's ~2-min post-provision hang (15.8.7). The
    networking mode is **gated on the LAN-P2P signal** ('.lan_p2p_detected', falling back to '
    TAILSCALE_ALLOW_LAN_P2P' in '.env'): a trusted home network requests '--network-address --network-
    mode bridged' (or 'shared' on WPA2-Enterprise), while the common AP-isolated / untrusted case boots
    plain fast user-mode networking ('colima start --memory 0.6 --vm-type vz').
106 - **Why Bridged Networking (when enabled)?** The '--network-mode bridged' and '--network-address'
    flags force the VM to receive its own IP directly from your local router's DHCP server, placing it
    on the same physical subnet as the host. This bypasses macOS network isolation and is what lets **
    external LAN devices** (phones, laptops on the same LAN), mDNS, and local service discovery reach
    across the container boundary. It requires the 'socket_vmnet' helper and is only worthwhile on a
    trusted LAN, which is why it is gated rather than always requested. **It is *not* required for
    direct hostcontainer P2P** that path (the gateway its own Mac) works even in user-mode networking
    via the routing-fix '.host_ips' RETURN-rule carve-out; see 15.3.5 and 15.8.7.
107 - **Enterprise Wi-Fi Fallback & Dynamic Roaming:** 'toggle.sh' dynamically queries 'system_profiler
    SPAirPortDataType' on boot. If it detects a **WPA2-Enterprise** connection (802.1X), it forces
    Colima to fall back to '--network-mode shared'. This prevents aggressive network intrusion systems
    on corporate/university Wi-Fi switches from terminating the host's Wi-Fi connection due to "MAC
    spoofing" (detecting multiple MAC addresses originating from a single authenticated port). P2P
    connections are impossible on these networks anyway due to AP Client Isolation.
108 - **Dynamic Roam Downgrading:** 'toggle.sh' writes its selected network mode to '/tmp/
    nullexit_colima_mode.txt'. When roaming between networks (e.g., Home to University), 'recover.sh'
    actively checks this state against the new network's security requirement ('.lan_p2p_detected'). If
    a mismatch is detected (e.g., bridged Colima on an Enterprise network), 'recover.sh' automatically
    forces a full 'toggle.sh --restart' in the background to safely transition the Virtual Machine's
    network mode and protect the host from de-authentication.
109 5. **Configure VM swap** 400MB swap file inside the VM to prevent OOM
110 6. **Clean corrupted AdGuard config** Remove empty 'AdGuardHome.yaml'
111 7. **Start containers** 'docker compose up -d'
112 8. **Wait for gateway Tailscale** Poll 'tailscale status' for "offers exit node" (up to 60s, abort at 40
    consecutive NoState)
113 9. **Resolve gateway IP** From '.gateway_ip' or 'docker compose exec tailscale tailscale ip -4'
114 10. **Connect host to mesh** Verify 'tailscaled' is running (auto-start if needed), then 'tailscale up
    --reset --ssh=true --accept-dns=false --accept-routes=true --exit-node=' ('--accept-routes=true'
    must be explicit '--reset' reverts it to 'false' otherwise, silently preventing the default route
    from switching to 'utun')
115 11. **Pre-flight checks** [1/3] 'tailscale ping' gateway, [2/3] 'dig +tcp AdGuard DNS, [3/3] WARP
    container internet
116 12. **If all pass** Set exit node + hijack DNS to gateway IP (single server, no fallback)
117 13. **If any fail** Enable SOCKS5 proxy + local DNS proxy as fallback
118 14. **Final DNS re-force** 'force_dns_to_gateway' called again after exit-node transition (tailscaled
    clobbers DNS briefly)
119 15. **Enable sleep prevention** 'caffeinate -i' in background to prevent idle sleep
120
121 ### STOP Path (gateway ON OFF)
122 1. Disconnect host Tailscale ('tailscale down')
123 2. 'docker compose down -t 5'
124 3. Reset DNS to 1.1.1.1

```



```

125 4. Stop local DNS proxy
126 5. **Stop sleep prevention** Kill caffeinate process
127 6. Full network state cleanup (proxies off, DNS cache flush, route flush, Wi-Fi power cycle)
128
129 ### Safety Mechanisms
130 - 'run_with_timeout()' Pure-bash watchdog for every system command (15s/120s)
131 - 'cleanup_handler()' Trap on ERR/INT/TERM/HUP restores DNS to 1.1.1.1, kills caffeinate, and tears down
    Tailscale
132 - 'force_dns_to_gateway()' Sets DNS and VERIFIES via 'networksetup -getdnsservers' (3 attempts with
    backoff)
133 - 'SUCCESS_RUN' flag Cleanup only runs if script didn't complete successfully
134 - 'start_sleep_prevention()' / 'stop_sleep_prevention()' Manages 'caffeinate -i' via PID file at '/tmp/
    nullexit-caffeinate.pid'
135
136 ---
137
138 ## 5. Routing & Firewall Architecture (Inside Container)
139
140 ### IP Rules ('ip rule show')
141 | Priority | Match | Table | Purpose |
142 |-----|-----|-----|-----|
143 | 99 | 'to 100.64.0.0/10' | 52 | CGNAT range Tailscale routes (injected by routing-fix) |
144 | 100 | 'from 172.18.0.2' | 200 | Container's own traffic (SOCKS5 proxy) tun0 |
145 | 101 | all | 51820 | Gluetun's WireGuard fwmark tun0 |
146
147 ### Table 200 (SOCKS5 proxy traffic)
148 - '162.159.192.1 via $DOCKER_GW dev eth0' WARP endpoint exception (prevents tunnel loop)
149 - 'default dev tun0' Everything else through WARP
150
151 ### iptables (TWO separate stacks!)
152 - **nftables backend** ('iptables') Used by Gluetun's 'post-rules.txt'. FORWARD policy DROP.
153 - **legacy backend** ('iptables-legacy') Used by Tailscale's 'ts-forward'. FORWARD policy ACCEPT.
154 - Both stacks apply to every packet. FORWARD RELATED,ESTABLISHED must exist in BOTH.
155
156 ### post-rules.txt (loaded by Gluetun)
157 '''
158 FORWARD: ACCEPT for tailscale0 in/out
159 INPUT: ACCEPT for tailscale0
160 NAT POSTROUTING: MASQUERADE on tun0
161 MANGLE: TCP MSS clamp to 1180
162 NAT PREROUTING: Redirect DNS (53) 5335 on tailscale0 and eth0
163 ip6tables: DROP FORWARD on tailscale0
164 '''
165
166 ---
167
168 ### 5.1 Tailscale Local Network Discovery (DERP Bypass)
169 By default, Gluetun's strict NAT swallows all of Tailscale's UDP hole-punching packets, forcing Tailscale
    to fall back to incredibly slow DERP relay servers (often adding 500ms+ latency).
170
171 To fix this, we use 'FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8' in 'docker-
    compose.yml' (on the 'warp' container). This creates 'ip rule' bypasses (Priority 99, Table 199)
    that route all packets destined for private IP ranges *outside* the WARP tunnel directly to the host
    's 'eth0'.
172 - **172.16.0.0/12** is especially critical because many modern Wi-Fi routers (and Docker itself) assign
    IPs in this block (e.g., '172.17.x.x').
173 - This bypass allows Tailscale's UDP packets to reach devices on the same Wi-Fi directly, establishing a
    fast peer-to-peer connection while the actual web traffic inside the tunnel remains fully routed
    through WARP.
174
175 ---
176
177 ### 5.2 Firewalling & Per-Device Access Control
178 Because every mesh device's traffic passes through the 'warp' container's network namespace, this
    namespace serves as the ultimate choke point for network-wide firewall rules.
179
180 **Where to Put Rules**
181 The right place to add firewall rules is 'scripts/routing-fix.sh'. It runs a 30-second loop inside the
    namespace and already handles re-injecting rules that Gluetun resets on reconnect. **Do not use '
    post-rules.txt' for dynamic rules**, as Gluetun flushes and resets it upon VPN reconnects.
182
183 **The Dual iptables Backend Constraint (CRITICAL)**
184 The container runs two simultaneous iptables backends: 'iptables' (for the nftables backend used by
    Gluetun) and 'iptables-legacy' (for Tailscale). Every rule **must** be added to both stacks, or it
    will silently fail for certain traffic.
185
186 **Avoiding Duplication**
187 Because 'scripts/routing-fix.sh' runs in a loop, you must always check for a rule's existence with '-C'
    before appending with '-A'. Otherwise, your rules will duplicate every 30 seconds.
188
189 **Per-Device Identification & Filtering**

```

```

190 Each mesh device has a stable '100.x.x.x' Tailscale IP, which is visible as the 'src' IP in the 'FORWARD'
    chain. This is how you identify devices for filtering.
191 * **Domain-level:** AdGuard Home (port 3000, credentials 'admin/nullexit') provides a REST API that
    supports per-client blocking rules based on their Tailscale IP.
192 * **IP-level:** Use 'iptables' in 'scripts/routing-fix.sh' (e.g., 'iptables -I FORWARD -s 100.x.x.x -d <
    blocked_ip> -j DROP').
193 * **Country-level:** Add 'ipset' to the 'routing-fix' apk install step in 'docker-compose.yml', and load
    CIDR ranges from 'ipdeny.com' into an ipset, then block that set in the 'FORWARD' chain. See 5.3 (
    Geo-IP Blocking) for the full configuration flow.
194
195 ### 5.3 Geo-IP Blocking
196 To block traffic to and from specific countries, the system dynamically downloads country-specific IP
    ranges from [ipdeny.com](http://www.ipdeny.com/ipblocks/data/countries/) and drops them via iptables
    'FORWARD' rules.
197
198 To add a new country to the blocklist:
199 1. Find the 2-letter ISO country code (e.g., 'cn' for China, 'ru' for Russia, 'il' for Israel, 'kp' for
    North Korea).
200 2. Open your '.env' file.
201 3. Update the 'BLOCKED_COUNTRIES' variable: e.g., 'BLOCKED_COUNTRIES="kp il cn ru"'.
202 4. Run 'toggle.sh' to restart the gateway and apply the new environment variables to the routing-fix
    container.
203
204 > [!NOTE]
205 > **Limitations of Geo-IP Blocking:** IP-based blocking is highly effective at blocking servers, direct
    apps, and raw infrastructure physically located and registered in a specific country (e.g., 'www.gov
    .il'). However, it will not block high-profile websites (e.g., 'jpost.com') that route their
    traffic through global Content Delivery Networks (CDNs) like Google Cloud, Cloudflare, or Fastly.
    Because the CDN's IP addresses are registered in the US or globally, the network traffic physically
    never routes to the blocked country, allowing it to bypass the Geo-IP firewall.
206
207 ## 6. macOS Quirks to Remember
208
209 1. Colima SSH tunnel is TCP-only UDP ports (like 5354/udp) are NOT accessible from host. Use 'dig +
    tcp' or the Python DNS proxy (UDPTCP converter).
210 2. 'docker compose exec -T' injects '\r' Always 'tr -d '\r'' when capturing output for comparisons
    or 'networksetup' commands.
211 3. 'brew services' can get stuck Fix with 'sudo brew services restart tailscale'. Common state: '
    brew services list' shows 'started' but 'tailscale status' shows "Tailscale is stopped."
212 4. No 'timeout' command macOS lacks GNU 'timeout'. Use the 'run_with_timeout()' bash function.
213 5. 'sudo -v' prompts even with NOPASSWD Use 'sudo -n' (non-interactive) everywhere.
214 6. DNS is scoped to en0 macOS scutil DNS resolver is scoped to en0 (Wi-Fi). DNS changes on other
    interfaces (like USB Ethernet) are ignored unless en0's service is also updated. The script always
    sets DNS on BOTH '$ACTIVE_SERVICE' and '$ENO_SERVICE'.
215 7. tailscaled clobbers DNS during exit-node transition 'force_dns_to_gateway' is called a second
    time at the end of the ENABLE branch to counteract this.
216 8. 'tailscale up --reset' resets all unspecified flags to defaults Every '--reset' call MUST also
    pass '--accept-dns=false' explicitly or tailscaled re-enables DNS management.
217 9. 'docker compose ps' CWD pitfall 'docker compose' commands require a 'docker-compose.yml' in the
    current working directory. Use 'docker ps' for raw container listing from any CWD; 'docker compose
    ps' requires the project directory.
218
219 ---
220
221 ## 7. Environment Variables
222
223 ### '.env' File
224 | Variable | Purpose |
225 |-----|-----|
226 | 'WIREGUARD_PRIVATE_KEY' | WARP WireGuard private key (from 'wgcf generate') |
227 | 'WIREGUARD_PUBLIC_KEY' | WARP WireGuard public key |
228 | 'WIREGUARD_ADDRESSES' | WARP WireGuard address (e.g., '172.16.0.2/32') |
229 | 'TS_AUTHKEY' | Tailscale auth key for the container |
230 | 'NULLEXIT_SEED' | 256-bit seed for HMAC-SHA256 script integrity signing (generated by 'setup.sh') |
231 | 'GATEWAY_RULE_PROFILE' | Rule compilation tier: 'light', 'medium', 'heavy' |
232 | 'GATEWAY_BYPASS_PING' | (Optional) 'true' to proceed even if pre-flight checks fail |
233 | 'GATEWAY_USE_EXIT_NODE' | (Optional) 'false' to skip exit node, DNS-only mode |
234 | 'GATEWAY_HIJACK_HOST' | (Optional) 'false' to skip DNS hijacking on the host (VPN/adblocking for
    Tailscale peers only) |
235 | 'GATEWAY_MSS' | (Optional) TCP MSS clamp value. Default and only validated-safe value is '1120' 15.3.2
    found 1180 (and 1160) still too large for the double-encapsulation overhead and causes mid-transfer
    stalls. Applied to both the container ('post-rules.txt') and the host ('pf.conf'), kept in sync. |
236 | 'WARP_FAIL_THRESHOLD' | (Optional) Consecutive 'warp=off' polls before auto-shutdown (default 6 = 30s;
    3 = 15s) |
237 | 'WARP_ENDPOINT_1' | (Optional) Override Cloudflare WARP WireGuard endpoint IP 1 (default
    '162.159.192.1') |
238 | 'WARP_ENDPOINT_2' | (Optional) Override Cloudflare WARP WireGuard endpoint IP 2 (default
    '162.159.193.1') |
239 | 'KILL_SWITCH' | (Optional) 'true' to enforce strict PF lock drops all host traffic if VPN fails.
    Breaks SSH if VPN dies. |

```

```

240 | 'STOP_COLIMA_ON_EXIT' | (Optional) 'true' to fully shut down the Colima VM on toggle-off (saves battery
    | on dedicated hosts) |
241 | 'ADGUARD_USER' | (Optional) AdGuard Home web UI username (default: 'admin') |
242 | 'ADGUARD_PASSWORD' | (Optional) Shared password for AdGuard Home and Tor ControlPort (default: '
    | nullexit') |
243 | 'BLOCKED_COUNTRIES' | Space-separated 2-letter ISO codes to block via ipdeny.com CIDR ranges (e.g. "kp
    | il cn ru") |
244
245 ---
246
247 ## 8. Diagnostic Commands
248
249 ### Container Health
250 ```bash
251 docker ps --format '{{.Names}} {{.Status}}'
252 docker compose exec -T tailscale tailscale status
253 docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace
254 ```
255
256 ### SOCKS5 Proxy
257 ```bash
258 curl --socks5-hostname 127.0.0.1:1080 --max-time 5 -s https://www.cloudflare.com/cdn-cgi/trace
259 ```
260
261 ### AdGuard DNS
262 ```bash
263 dig +tcp @127.0.0.1 -p 5354 google.com +short +timeout=5
264 ```
265
266 ### Firewall & Routing
267 ```bash
268 # nftables FORWARD chain
269 docker compose exec -T warp iptables -L FORWARD -v --line-numbers
270 # Legacy FORWARD chain
271 docker compose exec -T warp iptables-legacy -L FORWARD -v --line-numbers
272 # IP rules + routing tables
273 docker compose exec -T warp ip rule show
274 docker compose exec -T warp ip route show table 199
275 docker compose exec -T warp ip route show table 200
276 ```
277
278 ### Host macOS
279 ```bash
280 tailscale status
281 brew services list | grep tailscale
282 networksetup -getdnsservers Wi-Fi
283 networksetup -getsocksfirewallproxy Wi-Fi
284 sysctl net.inet.ip.forwarding
285 ```
286
287 ---
288
289 # Part II Design Decisions & Threat Model
290
291 ## 9. Privacy Architecture & Threat Model
292
293 ### Layered Defense
294 The gateway stacks four layers targeting different attack surfaces:
295
296 - **Layer 1 ISP-Level Surveillance (WARP):** Your ISP sees only an encrypted WireGuard tunnel to a
    | Cloudflare IP. In the US, ISPs can legally sell your browsing metadata and are subject to secret
    | NSLs (National Security Letters). WARP removes their visibility entirely.
297 - **Layer 2 DNS Tracking (AdGuard Home):** DNS is the phone book of the internet. By default, queries go
    | to your ISP's resolver, giving them a complete log of your traffic. AdGuard intercepts all DNS from
    | mesh devices, blocks trackers, and resolves queries through the WARP tunnel.
298 - **Layer 3 Device Identity & IP Exposure (Tailscale):** On untrusted networks (coffee shops, public Wi-
    | Fi), your device's IP is exposed. Tailscale creates an encrypted WireGuard mesh between your devices
    | , routing all traffic to your gateway exit node.
299 - **Layer 4 Kernel-Level IP Blocking (ipset/iptables):** Sophisticated malware bypasses DNS entirely
    | with hardcoded IPs. The 'rule-compiler' fetches ~16,700 threat-intelligence IPs/CIDRs (Spamhaus,
    | Feodo Tracker, Emerging Threats, CINS) and enforces them as 'DROP' rules in the kernel 'FORWARD'
    | chain blocking both outbound C2 connections and inbound attack traffic.
300
301 ### The Documented Threat: Why This Matters
302 This architecture is calibrated against documented mass-surveillance programs revealed in the Snowden
    | disclosures:
303 - **PRISM:** Bulk collection of communication content from major tech providers.
304 - **UPSTREAM:** Tapping physical backbone fiber, collecting metadata and content in bulk.
305 - **XKeyscore:** Retroactive search of unencrypted traffic.
306 - **MUSCULAR:** Infiltration of unencrypted internal datacenter interconnect links.
307

```

```

308 Passive bulk collection is not paranoia it is a documented reality. By double-encrypting and routing
    traffic, nullexit prevents your ISP from harvesting your metadata.
309
310 ### Trust Assumptions & Shifting
311 Every component shifts trust rather than eliminating it:
312 1. Physical Device Integrity: You trust your physical devices are not compromised.
313 2. Cloudflare Integrity: You trust Cloudflare not to correlate your WARP tunnel with exit traffic.
314 3. Tailscale Integrity: You trust Tailscale's coordination server not to MITM WireGuard keys.
315
316 The goal: no single provider has a complete picture. Your ISP sees an encrypted tunnel. WARP sees traffic
    with no account identity. Tailscale sees node topology but not content.
317
318 ### 9.1 Exit Node IP Rotation (Design Decision)
319
320 A common question for privacy architectures is whether the system should aggressively rotate its public
    exit IP (e.g., every 5 minutes, like a Tor circuit or a proxy rotator).
321
322 'nullexit' uses Cloudflare WARP via Gluetun. WARP provides anonymity in a crowd rather than
    aggressive rotation. Thousands of users share the same Cloudflare Edge IP simultaneously, making
    your traffic indistinguishable from the herd. The IP may occasionally rotate on its own (
    historically surfaced as a 'ROTATE' event in 'output.log'), but it remains generally stable.
323
324 Why aggressive IP rotation was intentionally excluded:
325 1. Broken TCP Connections: Every rotation instantly drops all long-lived TCP sessions (SSH, large
    downloads, WebSockets, gaming, Zoom).
326 2. CAPTCHA & 2FA Hell: Modern web security (e.g., Cloudflare, Akamai) aggressively flags IP-hopping
    as bot behavior. Aggressive rotation guarantees the user will be bombarded with "Verify you are
    human" CAPTCHAs and "New Login Detected" 2FA emails constantly.
327 3. WireGuard Architecture: WireGuard is stateless and does not natively support IP hopping on the fly
    . To rotate an IP, the VPN client must actively drop the tunnel, request a new endpoint, and re-
    handshake. This introduces latency gaps where the kill-switch would repeatedly trigger and blackhole
    the user's internet.
328
329 Verdict: For a daily-driver secure gateway, shared-IP crowding (WARP) is superior to aggressive
    rotation. As an added benefit, Cloudflare WARP IPs are actually highly static (tied to the
    persistent WireGuard public key generated for your device identity). Because your IP stays static
    and is part of Cloudflare's own trusted network, it builds a high "trust score," virtually
    eliminating CAPTCHA loops while still hiding you within the massive crowd of other users in that
    data center. All mesh devices routing through your gateway will reliably share this exact same
    trusted IP.
330
331 Note: The only exception where aggressive IP rotation is genuinely required is for specialized
    offensive security tasks like high-volume web scraping or dark web research where avoiding IP bans or
    active tracking overrides the need for TCP stability.)*
332
333 ### 9.2 Cloudflare WARP Privacy & Logging
334
335 Because 'nullexit' uses Cloudflare WARP (via Gluetun) as its upstream exit node, the system inherits
    Cloudflare's consumer privacy policy.
336
337 What Cloudflare DOES NOT log (Strict No-Logs Policy):
338 * Browsing History: They do not log the domains or URLs you visit.
339 * Traffic Destinations: They do not log the IP addresses of the servers you connect to.
340 * Content: They do not inspect or log the payload/data of your traffic.
341 * Data Brokering: They explicitly guarantee they will never sell or rent your data to advertisers.
342
343 What Cloudflare DOES log (Minimal Telemetry):
344 * Data Transfer Volume: Total bandwidth used (required to enforce 10GB/mo quotas if you are not using
    a paid WARP+ key).
345 * Performance Metrics: Anonymized connection speeds and latency to optimize their network routing.
346 * Aggregate Volume: Total traffic volumes to popular destinations in aggregate (e.g., "datacenter X
    sent 50TB to Netflix"), stripped of all user IDs.
347 * Device ID (Public Key): They store the persistent WireGuard public key generated by your container
    to authorize the connection.
348
349 Privacy Verdict:
350 Cloudflare WARP provides excellent historical privacy. Because they do not log your destinations,
    they cannot comply with historical subpoenas asking what websites you visited yesterday. However, it
    is not absolute anonymity. They still know your real ISP IP address while you are actively
    connected. For daily-driver privacy against ISPs, hackers, and corporate trackers, it is top-tier.
    For nation-state evasion, use Tor (see 11).
351
352 ### 9.3 OPSEC: Python Runtime Airgap (Isolated Mode)
353 'nullexit' explicitly relies on the host's native system Python 3 runtime for critical components like '
    scripts/dns-proxy.py'. Because this script runs with elevated privileges ('sudo -n') to bind to the
    privileged UDP/53 socket, modifying or polluting this Python environment is strictly forbidden.
354
355 The Zero-Dependency Rule & Isolated Mode ('-I')
356 You must NEVER install third-party 'pip' packages (e.g., 'requests', 'pyyaml') into the system Python
    environment that 'nullexit' uses.

```

```

357 - All Python scripts in this repository must be written using *only* the Python Standard Library ('socket
    ', 'urllib', 'json', 'ssl').
358 - Third-party packages introduce a massive supply-chain attack vector. A compromised 'pip' package
    installed globally could allow local malware to hook into the Python runtime and exploit the 'sudo'
    privilege of the DNS proxy to gain full root access.
359 - To enforce strict immunity against user-installed malicious packages, 'toggle.sh' explicitly invokes
    the proxy using **Python's Isolated Mode** ('python3 -I'). This flag forces the interpreter to
    completely ignore 'PYTHONPATH', 'PYTHONHOME', and the user's 'site-packages' directory, executing
    exclusively from the read-only, Apple-signed system framework.
360
361 ### 9.4 Recommended Upgrades
362
363 | Layer | Default | Upgraded |
364 |---|---|---|
365 | VPN Exit | Cloudflare WARP (US) | Mullvad (Swedish, proven no-logs, anonymous payment) |
366 | DNS Resolver | AdGuard via WARP | AdGuard via Mullvad DoH (Cloudflare-blind) |
367 | Coordination | Tailscale (US, OAuth) | Headscale (self-hosted, no third-party) |
368 | Auth | OAuth / persistent keys | Ephemeral auth keys (no identity persistence) |
369 | Content | WireGuard asymmetric | WireGuard + PSKs (coordination server blind) |
370
371 ##### Mullvad (Replacing WARP)
372 Swedish-based. Police raided their offices in 2023 and left empty-handed no logs exist to seize.
    Accounts are random numbers. Payment by cash or Monero. Replace 'VPN_SERVICE_PROVIDER=custom' with '
    VPN_SERVICE_PROVIDER=mullvad' in 'docker-compose.yml'.
373
374 ##### Mullvad DoH Upstreams
375 Point AdGuard's upstream resolvers to 'https://adblock.dns.mullvad.net/dns-query'. Cloudflare WARP only
    sees encrypted HTTPS to Mullvad's IPs completely blind to your DNS queries.
376
377 ##### Headscale
378 Self-hosted Tailscale coordination server. Eliminates Tailscale the company, keeping all node topology
    under your control.
379
380 ##### Ephemeral Auth Keys
381 Generate in the Tailscale admin panel. Used once to authenticate; node disappears from the admin panel
    after disconnect. Breaks persistent identity linkage.
382
383 ##### Pre-Shared Keys (PSKs)
384 Add a WireGuard PSK between peer nodes. Adds a symmetric encryption layer that Tailscale's coordination
    server has no knowledge of, blinding it from reading transit content.
385
386 ### 9.5 Structural Limitations
387 - **Traffic Analysis / Metadata:** Passive fiber tapping (UPSTREAM) can observe timestamps, data volumes,
    and IP ranges without reading content. Defeating traffic analysis requires mixing networks (Tor) or
    continuous dummy packet padding both heavily degrade performance.
388 - **Targeted State-Level Adversaries:** Consumer-grade VPNs are not a complete shield against a nation-
    state actively targeting you. This stack is designed to defeat bulk passive surveillance and
    corporate dragnet collection.
389
390 ---
391
392 ## 10. Per-Device Access Control
393
394 Because every mesh device routes internet-bound traffic through the 'warp' container's 'FORWARD' chain,
    they all appear with their stable '100.x.x.x' Tailscale IP as the 'src' address. This gives two
    powerful enforcement surfaces:
395
396 ### IP & Port Level ('scripts/routing-fix.sh')
397 Write raw 'iptables' rules targeting specific devices. The 30-second idempotent loop auto-survives
    Gluetun reconnects:
398
399 ```bash
400 # Block a specific device from a target subnet
401 iptables -I FORWARD -s 100.x.x.x -d 203.0.113.0/24 -j DROP
402
403 # Restrict a device to only HTTP/HTTPS
404 iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 3000 -j DROP
405 iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 443 -j DROP
406
407 # Block a device from internet egress entirely (mesh only)
408 iptables -I FORWARD -s 100.x.x.x -j DROP
409
410 # Time-based (e.g., cut off 10 PM 7 AM)
411 iptables -I FORWARD -s 100.x.x.x -m time --timestart 22:00 --timestop 07:00 -j DROP
412 ```
413
414 ### DNS Level (AdGuard REST API)
415 AdGuard Home supports per-client rules keyed by Tailscale IP:
416
417 ```bash
418 curl -X POST http://localhost:80/control/clients/add \

```

```

419 -H "Content-Type: application/json" \
420 -d '{
421     "name": "kids-ipad",
422     "ids": ["100.x.x.x"],
423     "blocked_services": ["youtube", "tiktok", "instagram"],
424     "safesearch_enabled": true
425 }',
426 ''
427 ---
428
429
430 ## 11. Tor & OPSEC Architecture
431
432 ### 11.1 Tor-over-WARP Topology
433
434 When extreme anonymity is required, integrating the Tor network into 'nullexit' provides a powerful Tor-
435 over-VPN topology:
436 * **The University/ISP** sees only an encrypted Cloudflare WireGuard tunnel, rendering them completely
437 blind to the fact that you are using Tor (effectively bypassing Enterprise firewall Tor blocks).
438 * **Cloudflare** sees an encrypted TCP stream heading to a Tor Guard Node. They cannot see your DNS
439 lookups, the destination website, or the content of your traffic.
440 * **The Tor Exit Node** sees your traffic, but not your real IP (only the Cloudflare IP).
441
442 ##### The "Transparent Proxy" Trap (What NOT to do)
443 It is tempting to build a system-wide "Transparent Tor Proxy" that forces all host traffic (e.g.,
444 standard Chrome/Safari browsers, background iCloud syncs) through Tor automatically. **This is a
445 massive OPSEC trap and destroys anonymity.**
446 Modern operating systems and browsers will instantly leak your identity through background syncing (e.g.,
447 Apple Mail, Google Drive) and browser fingerprinting (canvas fingerprints, WebRTC, screen
448 resolution). The Tor network protects your *IP*, but if your payload contains identifying cookies or
449 unique fingerprints, the Tor Exit Node (which could be run by malicious actors) will instantly tie
450 your session to your real identity.
451
452 ##### The nullexit Implementation: Isolated Procedural Proxy
453 To safely integrate Tor without triggering the transparent proxy trap, 'nullexit' implements an **
454 Isolated Tor SOCKS5 Proxy** with **Procedurally Generated Ports**:
455 1. **Isolated Container:** A lightweight 'tor' Docker container runs securely inside the 'warp' network
456 namespace. It does not hijack system traffic.
457 2. **Opt-In Usage:** You must consciously configure a highly-hardened browser (like the official Tor
458 Browser or a dedicated Firefox profile) to use the proxy. This prevents accidental background sync
459 leaks.
460 3. **Procedurally Generated Ports:** Rather than hardcoding the standard proxy ports (like
461 '127.0.0.1:9050' or '1080'), 'toggle.sh' randomizes all internal proxy ports on every boot into the
462 ephemeral range (10000-65000) and silently writes them to '/tmp/nullexit-ports.env'. This completely
463 defeats static fingerprinting by malware.
464 4. **Kernel-Level Collision Avoidance:** To prevent the randomizer from picking a port already in use by
465 a root daemon or Apple service (which would crash the gateway), the script directly queries the
466 macOS kernel socket table ('netstat -an -p tcp') to verify the port is absolutely free before
467 assigning it.
468 5. **The Honey-Port Tripwire:** While procedural ports defeat static fingerprinting, advanced malware
469 might attempt a full sequential port scan of 'localhost' to find the hidden proxy. To counter this,
470 'toggle.sh' spawns a "Honey-Port" (a fake random port with a netcat listener attached). If any local
471 application attempts to connect to the Honey-Port, it instantly logs a critical security alert to '
472 output.log' and fires a macOS desktop notification, serving as an early-warning Intrusion Detection
473 System (IDS) for local malware.
474
475 > **Threat Model note:** The Tor SOCKS proxy and Honey-Port Tripwire only bind to '127.0.0.1' (loopback).
476 Port randomization and the Honey-Port only defeat blind, generic port-scanners they do NOT protect
477 against a targeted local process that reads '/tmp/nullexit-ports.env', greps process memory, or
478 runs 'lsof -i'. The proper mitigation is a host-level loopback-filtering firewall (LuLu/Little
479 Snitch). See 15.9 (Threat Model: Local Proxy Discovery) for the full analysis and the rogue-binary
480 verification test.
481
482 ### 11.2 Hardening: Tor Bridges and obfs4 Pluggable Transports
483 For users operating under severe Deep Packet Inspection (DPI), 'nullexit' supports configuring Tor to
484 connect exclusively through private bridges using the 'obfs4' pluggable transport.
485
486 When 'TOR_USE_BRIDGES=true' is set:
487 - **What it does:** It hides the entry IP from being matched against the public Tor consensus list, and
488 disguises the traffic's protocol fingerprint from the WARP exit provider (Cloudflare).
489 - **What it does NOT do:** It does *not* add any extra layers of anonymity beyond what Tor already
490 provides. It is purely a censorship-circumvention and obfuscation mechanism.
491 - **Limitations:** 'obfs4' is highly effective against passive DPI, but it is not guaranteed to defeat
492 active-probing-resistant detection by a highly sophisticated state-level adversary.
493
494 > **Architecture Rule:** The 'nullexit' gateway deliberately does NOT bundle, hardcode, or automatically
495 fetch bridge lines. Bridge lines are highly sensitive and expire. Users must obtain their own bridge
496 lines from 'https://bridges.torproject.org' and manually populate the 'tor-bridges.txt' file. If
497 bridges are enabled but the file is missing or empty, the Tor container will intentionally crash and
498 fail to start rather than silently falling back to public guard relays.

```



```

463 **The Bootstrapping Paradox (Out-of-Band Key Exchange)**
464 Automating the bridge acquisition process (e.g., scripting the gateway to log into the Telegram ‘
    @GetBridgesBot’ via MTPROTO API) introduces catastrophic OPSEC flaws:
465 1. **Identity Leaks:** Baking a Telegram session into the container explicitly links the "anonymous"
    gateway to a real-world physical phone number.
466 2. **The Chicken & Egg Deadlock:** In heavily censored regimes, Telegram itself is usually blocked. If
    the container requires Telegram to fetch bridges to bypass the firewall, it will fail because it
    cannot bypass the firewall to reach Telegram.
467 3. **Bridge Exhaustion & CAPTCHAs:** Automated scraping behaves identically to state-sponsored scrapers,
    triggering Tor’s anti-bot CAPTCHAs which leads to silent proxy failures and exhausts the limited
    public bridge pool.
468
469 Instead, users should rely on a less secure layer to manually reach Telegram, obtain the bridges, and
    populate ‘tor-bridges.txt’. This can be done via the standard ‘nulloxit’ WARP tunnel, a mobile phone
    on cellular data, Mullvad VPN, or a secure remote VPS. *(Note: We plan on fully integrating Mullvad
    and VPS architectures directly into ‘nulloxit’ in the future, but have not had the time to build it
    yet).* This intentionally "air-gaps" the cryptographic key exchange from the proxy infrastructure,
    leaving zero API tokens and zero network traces for an adversary to exploit.
470
471 ### 11.3 Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification
472
473 #### Context
474 Unlike the host’s SOCKS5 fallback path (whose kill-switch interaction is documented in 15.7), the ‘tor’
    container’s connection status and fail-closed behavior under VPN tunnel drops is not governed by
    macOS ‘pf’ rules. Instead, it relies on network namespace inheritance inside Docker Compose.
475
476 #### Mechanics
477 1. **Shared Network Namespace:** The ‘tor’ service is configured with ‘network_mode: service:warp’ in ‘
    docker-compose.yml’. This shares the network namespace of the ‘warp’ (Gluetun) container.
478 2. **Inherited Egress Rules:** All socket communication within the namespace is bound to the same
    networking stack. Gluetun manages this stack by redirecting traffic through ‘tun0’ (the WireGuard/
    WARP tunnel) and applying ‘iptables’ rules that explicitly drop outbound traffic originating on ‘
    eth0’ (the physical/bridge interface), except for permitted local subnets or the VPN server’s
    endpoint.
479 3. **Tunnel Fail-Closed:** If the WARP connection drops or the ‘tun0’ interface is brought down, the
    routing table entries for ‘tun0’ become invalid or disappear. Because the ‘iptables’ rules in the
    shared namespace continue to drop any outgoing internet traffic attempting to exit via ‘eth0’, the ‘
    tor’ container cannot leak raw traffic to the host’s underlying networks. Egress fails closed.
480
481 #### Verification / Manual Test Case
482 To manually verify that the Tor container fails closed and does not leak traffic when the tunnel dies:
483
484 1. **Verify active path:** Ensure the gateway is active (‘toggle.sh start’) and fetch the Tor port (‘
    $TOR SOCKS_PORT’ in ‘/tmp/nulloxit-ports.env’). Verify Tor traffic routes correctly:
485     ‘‘‘bash
486     curl --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.org/api/ip
487     ‘‘‘
488     *(This should output a Tor exit node IP).*
489
490 2. **Simulate tunnel drop:** Bring down the virtual interface inside the shared network namespace:
491     ‘‘‘bash
492     docker compose exec warp ip link set dev tun0 down
493     ‘‘‘
494
495 3. **Re-run the egress check:**
496     ‘‘‘bash
497     curl --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.org/api/ip
498     ‘‘‘
499     **Expected Result:** The request must hang and eventually time out, or immediately fail with a proxy
    error (e.g., ‘curl: (97) SOCKS5: connection failed’ or ‘curl: (7) Failed to connect’).
500     **Leaked Egress Check:** Check the host’s Wi-Fi router or firewall logs. No outbound packets from
    the Tor container should route directly over the bridge network to the internet.
501
502 ### 11.4 Transparent .onion Routing via RFC 2544 Subnet (198.18.0.0/15)
503
504 #### Context
505 To enable seamless dark web browsing across the entire mesh, the gateway requires a mechanism to
    dynamically map ‘.onion’ domains to IP addresses that the host operating system and network layers
    can route.
506
507 #### IP Address Conflict & Solution
508 1. **Initial Conflict:** Initially, the Tor virtual address network was configured to map ‘.onion’ sites
    within the ‘10.192.0.0/10’ range. However, the Colima Docker daemon dynamically allocated
    ‘10.200.1.0/24’ to the containers. Because this Docker subnet mathematically falls within the
    ‘10.192.0.0/10’ block, a routing loop occurred: all host packets destined for the Docker containers
    on ports like ‘9050’ or ‘9051’ were intercepted by the transparent proxy NAT rules and dropped,
    causing proxy failures.
509 2. **Tailscale Private IP Exclusion:** Shifting the subnet to another private range like ‘10.123.0.0/16’
    resolved the port conflict, but led to connection timeouts. On macOS, Tailscale’s Exit Node routing
    actively excludes RFC 1918 private subnets (e.g. ‘10.0.0.0/8’, ‘172.16.0.0/12’, ‘192.168.0.0/16’) to
    maintain accessibility to local LAN resources (like printers/routers). Thus, packets targeting

```



```

10.123.x.x' were dropped locally by Tailscale before reaching the gateway tunnel.
510 3. **The Benchmarking Subnet Solution:** To bypass these private IP exclusions without manual static
routing tables, Tor's virtual network was changed to **198.18.0.0/15** (reserved by RFC 2544 for
benchmarking). Since this is not a private RFC 1918 range, Tailscale exit-node routing automatically
encapsulates and routes all traffic destined for '198.18.0.0/15' through the tunnel to the gateway.
511 4. **Transparent NAT Redirection:** The 'routing-fix' sidecar applies NAT rules in the shared network
namespace:
512 'bash
513 iptables -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040
514 ''
515 This intercepts all incoming TCP traffic targeting the mapped '.onion' range and transparently proxies
it through Tor's transparent port ('9040').
516
517 #### Exit Node Exclusions
518 The gateway supports the ability to exclude specific countries from being used as exit nodes (e.g. to
avoid certain jurisdictions or nodes with poor network latency). This is configured globally via the
'TOR_EXCLUDE_EXIT_NODES' environment variable in '.env', which gets dynamically appended to Tor's '
ExcludeExitNodes' and strictly enforced with 'StrictNodes 1'.
519
520 ---
521
522 ## 12. Censorship-Resistant Transport & Egress Compartmentalization
523
524 ### 12.1 Why this is needed
525 WARP can be blocked from deployment countries with strict internet controls. The most likely cause isn't
that "encrypted traffic is blocked" in general it's that WireGuard has a recognizable handshake
signature, and 'WIREGUARD_ENDPOINT_IP=162.159.192.1' on 'WIREGUARD_ENDPOINT_PORT=2408' is a fixed,
easily-enumerable target (Cloudflare's known WARP range). That combination is exactly what IP/ASN-
based and DPI-based blocking is good at catching.
526
527 ### 12.2 Important correction: Gluetun's 'SHADOWSOCKS=on' is the wrong direction
528 Gluetun's built-in Shadowsocks option runs a Shadowsocks **server** inside the already-connected VPN
tunnel, so other LAN devices can reach out through Gluetun's *existing* connection via the SS
protocol. It is not a way to make Gluetun itself connect **outbound** through a Shadowsocks server
Gluetun only speaks OpenVPN or WireGuard as its actual transport. There is no 'VPN_TYPE=shadowsocks'
'. Confirmed against Gluetun's own maintainer/community discussion: people have asked for a "
Shadowsocks-client" mode specifically to chain through Gluetun, and the standing answer is "run a
separate Shadowsocks client container."
529
530 ### 12.3 Diagnose before building anything
531 The right fix depends entirely on what's actually being blocked:
532 - Point Gluetun at a **different provider/IP/ASN** (any other Gluetun-supported WireGuard/OpenVPN
provider, or a self-hosted WireGuard server) and test. If that gets through, the block is Cloudflare
/WARP-specific, not a general WireGuard block **no new component needed**, just swap '
VPN_SERVICE_PROVIDER' / the endpoint.
533 - If WireGuard fails regardless of endpoint, the block is protocol-level (DPI fingerprinting the
handshake) proceed to obfuscation.
534
535 ### 12.4 Path A: Swap WARP for a different Gluetun-native transport
536 Simplest fix, only applies if 12.3 shows the block is Cloudflare-specific. Change 'VPN_SERVICE_PROVIDER'
/ 'VPN_TYPE' / endpoint in 'docker-compose.yml' and in the now-centralized 'WARP_ENDPOINT_*' values
in '.env'. Nothing else in the stack needs to change.
537
538 ### 12.5 Path B: Add an obfuscation hop
539 Needed if WireGuard itself is being fingerprinted/blocked. Two ways to slot it in:
540
541 **B1 Wrap (recommended):** Run a Shadowsocks (or obfs4/v2ray) client as a new sidecar container. Point
Gluetun's 'WIREGUARD_ENDPOINT_IP' at '127.0.0.1:<local-port>' instead of Cloudflare directly; the
sidecar relays that traffic to a Shadowsocks server you run yourself abroad. Keeps Gluetun's kill-
switch, healthchecks, and the entire rest of the stack (Tailscale, AdGuard, ipset, 'routing-fix.sh')
untouched, since none of them know the transport changed underneath 'warp'.
542
543 **B2 Replace:** Swap the 'warp' service entirely for a Shadowsocks-client + 'tun2socks' container that
owns the network namespace instead of Gluetun. More invasive loses Gluetun's built-in kill-switch/
healthcheck, which would need to be rebuilt by hand (see Section 12.9 for the pluggable architecture
to support this natively).
544
545 **Recommendation: B1.** Smaller surface area, preserves the self-healing architecture already built and
documented.
546
547 ### 12.6 Fingerprinting caveat
548 Plain Shadowsocks can still be caught by active-probing DPI in more aggressive censorship environments.
If plain SS gets blocked too, the next step up is an obfuscation plugin ('v2ray-plugin', 'Cloak') or
a newer protocol designed against active probing (VLESS+Reality, Hysteria2). Test plain SS first
don't over-build before confirming it's needed.
549
550 ### 12.7 AmneziaWG (The High-Speed Alternative)
551 If you want to maintain the bare-metal speeds of WireGuard while bypassing DPI (e.g. in Egypt or Russia),
the best alternative to Shadowsocks/V2Ray is AmneziaWG. AmneziaWG is a fork of WireGuard that pads
the handshake packets with randomized "junk" data, entirely destroying the 148-byte DPI signature
that firewalls look for.

```

```

552 - **Pros:** Because it runs entirely as UDP without TCP-over-TCP overhead, it completely avoids "TCP
      meltdown" latency spikes associated with Shadowsocks/V2Ray wrappers.
553 - **Cons:** You cannot use Cloudflare WARP. Furthermore, it requires completely ripping Gluetun out of
      the 'nullexit' stack and building a custom Docker container, meaning you lose Gluetun's built-in
      kill-switches and health-check logic.
554
555 ### 12.8 Infrastructure Costs & Privacy
556 To use either Shadowsocks (Path B) or AmneziaWG, you cannot use the free Cloudflare WARP infrastructure,
      because Cloudflare servers only speak standard WireGuard. The software for both custom protocols is
      100% free and open-source, but you must host the remote server infrastructure yourself.
557 - **Free Tier (Oracle Cloud):** You can host your remote server on Oracle Cloud's "Always Free" tier for
      $0. However, this routes your highly sensitive, censorship-evading traffic directly through Oracle's
      massive corporate data broker. If privacy is the core objective, handing all your metadata to Oracle
      defeats the purpose.
558 - **Paid Tier (Hetzner / Mullvad):** The recommended path is to rent a standard ~$4/month Virtual Private
      Server (VPS) in a free country (e.g., Hetzner in Germany, subject to strict EU GDPR privacy laws)
      and self-host the server. Alternatively, Mullvad VPN (5/month) provides native v2ray/Shadowsocks
      bridges built into their network, allowing you to bypass censorship with a strict zero-logs policy.
      It is highly recommended to pay with untrackable payment methods to maintain complete anonymity,
      which these providers typically support without requiring local residency.
559
560 ### 12.9 The Future: Egress Compartmentalization (Pluggable Architecture)
561 To natively support invasive replacement paths like **Path B2** or **AmneziaWG** (Section 12.7) without
      breaking the core 'nullexit' routing and monitoring logic, the architecture must be refactored into
      a modular, "pluggable" design.
562
563 1. **The Egress Plugin System:** Currently, WARP-specific logic (like 'cdn-cgi/trace' health checks and
      hardcoded interface names) is scattered across 'toggle.sh' and 'scripts/routing-fix.sh'. This should
      be abstracted into a 'scripts/plugins/' directory, where each egress method ('warp.sh', 'v2ray.sh')
      implements standardized functions: 'get_egress_interface()', 'check_health()', and 'get_bypass_ips()'
564 2. **Dynamic Docker Compose:** Based on an 'EGRESS_TYPE' variable in '.env', dynamic compose file
      generation would spin up only the required egress containers (e.g., skipping the 'warp' container
      when 'EGRESS_TYPE=v2ray' is set).
565 3. **Agnostic Routing & Watchers:** 'routing-fix.sh' would dynamically read the 'EGRESS_INTERFACE' from
      the active plugin to apply 'iptables' rules, and the background watchers in 'toggle.sh' would become
      an agnostic 'Egress Watcher' that invokes 'check_health()' periodically.
566
567 This compartmentalization transforms 'nullexit' into a universal, censorship-resistant gateway where the
      entire egress layer can be swapped by changing a single '.env' variable and running './toggle.sh --
      restart'.
568
569 ---
570
571 ## 13. Roaming & Recovery Design: How nullexit Mimics Commercial VPNs
572
573 The 'nullexit' roaming/network recovery subsystem replicates the exact engineering mechanics used by
      commercial, premium VPN clients (like Mullvad or NordVPN) to transition across Wi-Fi networks safely
      and seamlessly. This is the design overview; the specific bugs encountered and fixed along the way
      live in the Incident Log (15.5 Roaming/Sleep-Wake, 15.7 Tailscale P2P/SNAT).
574
575 ### The 4-Step Transition Cycle
576 When your Mac disconnects from one Wi-Fi and associates with another, 'nullexit' coordinates the
      following events:
577
578 1. **Firewall Lock (Kill-Switch):** The macOS Packet Filter (PF) anchor 'com.apple/nullexit' remains
      locked on the physical interface ('en0'), blocking all direct outbound IPv4 and IPv6 traffic. This
      guarantees that your real ISP IP or unencrypted DNS requests **never leak** onto the new network
      while the connection is rebuilding.
579 2. **System Event Interception:** A launchd LaunchAgent ('watcher.sh') listens to the macOS System
      Configuration framework for the 'State:/Network/Global/IPv4' key changes, spawning 'recover.sh --
      post-wake' in under 500ms when a change is detected.
580 3. **Bypass Routing:** The script dynamically resolves the new physical network's IP gateway (e.g.
      '192.168.137.1'), cleans up stale bypass routes from the previous network, and adds static bypass
      routes pointing to Cloudflare's WARP endpoints and the Tailscale control plane directly via the new
      gateway. This allows the VPN client to negotiate its handshake outside of the tunnel.
581 4. **Hole-Punching / NAT Re-negotiation:** The script runs a target check on the 'warp' container's
      tunnel. If the WireGuard connection is wedged by the network switch, it force-recreates the
      container to establish a fresh NAT binding to Cloudflare, restoring connection in ~15 seconds.
582
583 ---
584
585 ## 14. Deployment, Packaging & Known Limitations
586
587 ### 14.1 Packaging nullexit as a macOS Application (.dmg)
588
589 nullexit can be packaged into a native '.app' bundle, but this introduces nested virtualization
      requirements.
590
591 nullexit uses **Colima** (Apple Virtualization Framework 'vz') to run a Linux VM hosting Docker.
      Installing the '.app' inside a virtualized macOS environment (Parallels, cloud Mac) means running a

```

```

Linux VM inside a macOS VM requiring hardware-level **nested virtualization**.
592
593 ##### Hardware Support
594 - **Supported:** Apple Silicon M3/M4 (macOS Sequoia 15+), Intel Macs with VMX passthrough.
595 - **Unsupported:** M1, M2, A18 Pro. Colima crashes silently; the terminal ('setup.sh') surfaces crash
    logs.
596
597 ##### Build Steps
598 ```bash
599 # 1. Verify nested virtualization support
600 sysctl -a | grep hv_nested_virt_supported # expect: 1
601
602 # 2. Compile the AppleScript launcher
603 osacompile -o "Nullexit.app" "Toggle Gateway.applescript"
604
605 # 3. Embed scripts into app resources
606 mkdir -p "Nullexit.app/Contents/Resources/scripts"
607 cp toggle.sh setup.sh recover.sh docker-compose.yml "Nullexit.app/Contents/Resources/"
608
609 # 4. Package into DMG
610 hdiutil create -volname "Nullexit Installer" -srcfolder "./Nullexit.app" -ov -format UDZO Nullexit.dmg
611
612 # 5. Bypass Gatekeeper (unsigned app)
613 xattr -cr /Applications/Nullexit.app
614 ```
615
616 Without a paid Apple Developer certificate ($99/yr), users see an "App is damaged" Gatekeeper warning and
    need to run step 5.
617
618 ### 14.2 Moonlight/Sunshine + Tailscale Race Condition
619
620 When streaming from a remote Windows host running Sunshine over a Tailscale mesh (e.g., while the client
    is on a cellular hotspot), a race condition can occur on boot:
621
622 1. Both Tailscale and Sunshine are set to start automatically on boot.
623 2. Tailscale takes a few seconds to fully initialize its '100.x.x.x' virtual adapter and establish a
    connection.
624 3. **Sunshine starts too fast.** It launches before Tailscale is fully ready. Since Sunshine only binds
    to network interfaces that are active at the exact moment of startup, it misses the Tailscale
    adapter entirely.
625 4. Moonlight clients (even when configured with the correct Tailscale IP) will fail to connect because
    Sunshine isn't listening on that interface.
626
627 **The "Magic Toggle" Symptom:**
628 If the user manually enables or disables Wi-Fi on the host PC, it triggers a global Windows network
    change event. Sunshine listens for these events, re-enumerates all network adapters, and says, "Oh,
    there's a Tailscale adapter here!" and finally binds to it. Suddenly, Moonlight can connect.
629
630 **The Permanent Fix:**
631 On the Windows host, open 'services.msc', locate the **Sunshine Service**, and change its Startup type
    from **Automatic** to **Automatic (Delayed Start)**. This forces Windows to wait a minute or two
    after booting before launching Sunshine, ensuring Tailscale's virtual adapter is fully initialized
    first.
632
633 ### 14.3 Chrome Remote Desktop Performance (UDP NAT)
634
635 ##### Observation (July 11, 2026)
636 When the user connects to their host Mac using Chrome Remote Desktop (CRD) while 'nullexit' is active,
    the connection suffers from massive latency, lag, and degraded video quality. AdGuard Home's '
    blocked.log' shows zero blocked DNS requests for Google, WebRTC, or 'chromoting' endpoints,
    indicating this is not a DNS sinkhole issue. *(See also 15.11 for the separate CRD-on-remote-device
    connection-failure quirk.)*
637
638 ##### Root Cause
639 This is a fundamental limitation of tunneling all host traffic through a strict commercial NAT (
    Cloudflare WARP).
640 1. Chrome Remote Desktop attempts to use **STUN (UDP hole-punching)** to establish a lightning-fast,
    direct peer-to-peer WebRTC connection between the client device and the host Mac.
641 2. Because all non-Tailscale traffic is forcefully routed through the 'warp' container, the STUN packets
    exit the network from a Cloudflare IP.
642 3. Cloudflare's enterprise NAT infrastructure strictly drops unsolicited inbound UDP packets. The UDP
    hole-punch fails.
643 4. Because a direct P2P connection cannot be established, CRD falls back to **Relay Mode** (TURN),
    bouncing all video data back and forth through Google's cloud servers instead of sending it directly
    over the local or mesh network. This relaying introduces massive latency.
644
645 ##### Workaround / Fix
646 This lag is the accepted "tax" for maintaining absolute cryptographic anonymity; the only way to restore
    CRD speed would be to leak its UDP traffic outside the tunnel (violating zero-trust).
647 To achieve low-latency remote access, users should bypass CRD entirely and use macOS's built-in **Screen
    Sharing (VNC)** directly over the Tailscale IP. Tailscale's sophisticated custom DERP/WireGuard

```

```

architecture successfully UDP hole-punches through strict NATs, providing a direct, encrypted, high-
speed connection without relying on external cloud relays.
648
649 ### 14.4 Dynamic IP Fetch vs. MagicDNS
650
651 ##### Observation
652 The 'nullexit' gateway uses a dynamically fetched IP address (cached in '.gateway_ip') to configure host
routing and the 'pf' kill-switch, rather than utilizing Tailscale's user-friendly MagicDNS hostname
(e.g., 'nullexit-gateway').
653
654 ##### Why MagicDNS is Not Used
655 1. **The "Chicken and Egg" Bootstrapping Problem:**
656 macOS's Packet Filter ('pf') and low-level routing commands ('route add') operate strictly at Layer 3
and require raw IP addresses. If 'toggle.sh' or 'recover.sh' attempted to use 'nullexit-gateway',
the host Mac would need to perform a DNS lookup to resolve it. However, the very first step of the
gateway's initialization is to completely hijack the host's DNS and point it *at the gateway itself
*. The Mac cannot resolve the gateway's MagicDNS name if it needs the gateway's IP to know where to
send the DNS query!
657 2. **Dynamic Self-Healing (Avoiding Stale Cache):**
658 To solve the bootstrapping problem without causing bugs if the node is deleted and re-authenticated on
the Tailscale Admin Console, 'toggle.sh' explicitly runs 'docker compose exec ... tailscale ip -4'
during boot to fetch the absolute freshest IP dynamically. It saves this to '.gateway_ip' so that
fast-roaming scripts like 'recover.sh' can instantly retrieve the routing target without incurring
the latency of querying Docker.
659
660 ### 14.5 Battery & Power Measurement Quirks
661
662 When trying to optimize this framework (or any daemon) for battery life, developers often look for a way
to programmatically measure the exact Wattage consumed by a specific process (e.g., "How many Watts
is 'toggle.sh' using?").
663
664 **This is a hardware impossibility on macOS and Linux.**
665
666 ##### Why Activity Monitor is Lying to You
667 Your machine's logic board only has physical power sensors for the *entire* CPU package, the GPU, and the
RAM. It knows the CPU is pulling 5W total, but it has no physical hardware capability to know
whether Docker is using 3W and Chrome is using 2W.
668
669 The "Energy Impact" score you see in macOS Activity Monitor is a **synthetic, heuristic score**
calculated by 'powerd'. It penalizes apps for waking up the CPU (idle wakes), doing disk I/O, or
keeping the screen awake. It is a relative score, not actual Wattage.
670
671 ##### The Proper A/B Testing Methodology
672 If you want to truly know how many Watt-hours or mAh your background daemons ('routing-fix.sh', the DNS/
WARP watchers, etc.) are costing you, you must perform a hardware-level A/B drain test:
673
674 1. Unplug the laptop, run the gateway, and note your exact battery mAh using:
675     'bash
676     ioreg -l | grep "AppleRawCurrentCapacity"
677     '
678 2. Leave the laptop idle for exactly 1 hour.
679 3. Run the command again to see exactly how many mAh were drained.
680 4. Turn the gateway off and repeat the test for another hour.
681
682 The delta in mAh drained between the two tests is the true, exact cost of running the software. When
optimizing polling loops (e.g., changing 'sleep 5' to 'sleep 30'), this is the only reliable way to
measure the actual battery life returned to the user.
683
684 ### 14.6 Host Lockdown Mode (Why It's Impossible on macOS)
685
686 A common privacy feature request is a "Host Lockdown Mode" (Mode 3): A mode where the Docker exit node
remains perfectly functional for remote devices, but the host machine itself is completely blocked
from accessing the internet (to prevent even encrypted host traffic from leaking metadata).
687
688 **Why this is impossible on macOS:**
689 On Linux (e.g., a Raspberry Pi), Docker runs natively. The host's traffic traverses the 'OUTPUT' iptables
chain, while Docker traffic traverses the 'FORWARD' chain. We could easily drop all 'OUTPUT'
traffic to kill the host's internet, and Docker would continue routing traffic perfectly.
690
691 However, Docker on macOS runs inside a Linux Virtual Machine (via 'qemu' or Apple 'Virtualization.
framework'). This VM runs as a standard macOS user-space application. If you configure the macOS
firewall ('pf') to drop all host outbound traffic, it will indiscriminately drop the VM application's
traffic as well, instantly killing the Cloudflare WARP tunnel and the exit node.
692
693 Writing complex macOS 'pf' rules to "allow the VM app, allow Cloudflare WARP IPs, allow Tailscale DERP
IPs, but block everything else" is highly fragile and heavily discouraged. Since the current 'toggle
.sh' default mode already fully encrypts the host's traffic via the WARP tunnel, the host is already
protected from ISP leaks.
694
695 ---
696

```

```

697 # Part III Incident Log & Resolved Issues
698
699 This is the single, deduplicated log of every incident, bug, and resolved issue. Entries are organized
    into subsystem groups (15.115.12). Each entry follows a **Symptom / Root Cause / Fix** shape where
    applicable; packet-level analyses are preserved as clearly-labeled **Deep dive** blocks. The terse
    dated changelog lives in 17; this section holds the full post-mortems.
700
701 ## 15. Incident Log
702
703 ### 15.1 Exit-Node Return Path & SOCKS5 Failover
704
705 #### 15.1.1 Exit Node Return Path ('rx 0') & The SOCKS5 Failover RESOLVED
706 **Symptom:** Gateway showed 'tx 936 rx 0' transmitted but never received return traffic. For days,
    Tailscale's exit node refused to return traffic back to the client, leading to the assumption that
    Tailscale's userspace forwarding was bypassing 'iptables' and ignoring the WARP 'tun0' interface. In
    a desperate attempt to fix this, a **SOCKS5 proxy** ('scripts/socks5-proxy.py') was introduced as a
    replacement.
707
708 **Root Cause:** Tailscale's userspace forwarding was not the problem. 'docker-compose.yml' included '
    FIREWALL_OUTBOUND_SUBNETS=100.64.0.0/10'. Gluetun created a strict routing rule ('ip rule add to
    100.64.0.0/10 lookup 199', priority 99) that forced all Tailscale CGNAT traffic to bypass the VPN
    via the unencrypted Docker bridge ('eth0'). Return packets were sent to the macOS host instead of
    back through 'tailscale0', blackholing all return packets back to the client.
709
710 **Fix:** Removed 'FIREWALL_OUTBOUND_SUBNETS' entirely, immediately restoring the Tailscale exit node. '
    scripts/routing-fix.sh' now injects 'lookup 52', and return packets correctly flow back into '
    tailscale0'. Instead of removing the SOCKS5 proxy, we repurposed it into a **bulletproof failover**
    for when restrictive Wi-Fi networks block Tailscale UDP ports.
711
712 Architecture after the fix:
713 ```
714 PRIMARY (Exit Node):
715 DNS/TCP/UDP: Apps Tailscale Exit Node gateway container tun0 WARP internet
716
717 FAILOVER (SOCKS5 - if Tailscale is blocked by local Wi-Fi):
718 DNS: Browser 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP
719 TCP: Apps SOCKS5:1080 gateway container tun0 WARP internet
720 Mesh: SSH/ping 100.x.x.x tun5 WireGuard peers (independent of proxy)
721 ```
722
723 **Deep dive: Exit Node Return Path Analysis (June 26, 2026).**
724 This preserves the detailed packet-level analysis that led to identifying and fixing the 'rx 0' bug.
725
726 *The exit node was probably never going to work in this container topology.* Here's the packet path
    traced by reading 'ip rule show' and the routing tables live:
727
728 **Outbound (Phone-1 internet):**
729 1. Phone-1 sends packet to gateway via Tailscale mesh
730 2. Gateway's tailscale0 (kernel TUN, 'TS_USERSPACE=false') decrypts packet enters FORWARD chain
731 3. Packet has src='100.87.42.87' (Phone-1), dst=some internet IP
732 4. FORWARD chain (nftables): Rule 1 matches ('-i tailscale0 -j ACCEPT')
733 5. Routing decision for the forwarded packet. Which ip rule matches?
734 - Rule 98: 'to 172.18.0.0/16' no (dst is internet)
735 - Rule 99: 'to 100.64.0.0/10' no (dst is internet)
736 - Rule 100: 'from 172.18.0.2' **NO** (src is '100.87.42.87', not the container IP!)
737 - Rule 101: 'not fwmark 0xca6c lookup 51820' **YES** (no fwmark on forwarded packets)
738 6. Table 51820: 'default dev tun0' packet goes out through WARP
739 7. NAT POSTROUTING: 'MASQUERADE on tun0' src becomes WARP IP
740 8. Packet exits through WARP to internet (in theory)
741
742 **Return (internet Phone-1):**
743 1. Return packet arrives on tun0, conntrack de-MASQUERADES dst back to '100.87.42.87'
744 2. Routing decision for dst='100.87.42.87':
745 - Rule 99: 'to 100.64.0.0/10 lookup 199' **YES**
746 3. Table 199: '100.64.0.0/10 via 172.18.0.1 dev eth0' **sends it to the Docker gateway!**
747 4. **This is wrong.** The packet should go to tailscale0 (to be re-encrypted and sent to Phone-1), but
    table 199 sends it out eth0 to the Docker bridge host.
748
749 **Table 199 is the smoking gun.** It has one route: '100.64.0.0/10 via 172.18.0.1 dev eth0'. This sends
    ALL Tailscale CGNAT return traffic out the Docker bridge to the host not back through tailscale0.
    The return packet leaves the container's network namespace entirely, hits the host's routing stack,
    and gets dropped because the host has no idea what to do with a raw packet for '100.87.42.87'.
750
751 **Why table 199 exists and why it's wrong for exit node traffic.* Table 199 + the CGNAT rule (pref 99) was
    designed so that the **container itself** can reach other Tailscale peers (e.g., for mesh
    management, DERP relay). For that use case, routing CGNAT traffic via the Docker bridge to the host
    (which has its own tailscaled) makes sense. But for **forwarded exit node traffic**, the return
    packet needs to go back through tailscale0 inside the container not out to the host. The CGNAT rule
    intercepts the return packet before it can be routed back to tailscale0 through the main table or
    Tailscale's own routing (table 52, which is at pref 5270).
752

```

```

753 *The SOCKS5 path works because it dodges all of this.* SOCKS5 proxy creates connections **from the
    container's own IP** ('172.18.0.2'). So:
754 - ip rule 100 ('from 172.18.0.2 lookup 200') matches
755 - Table 200 has 'default dev tun0'
756 - Return traffic comes back to '172.18.0.2' on tun0, conntrack de-NATs, proxy sends response to client
757
758 No FORWARD chain involved. No CGNAT rule involved. No table 199. It just works.
759
760 *Summary of findings:*
761
762 | Finding | Status | Evidence |
763 |-----|-----|-----|
764 | Table 199 CGNAT route hijacks exit node return traffic | **Fixed** | Changed 'lookup 199' to 'lookup
    52' in 'scripts/routing-fix.sh'. Return packets now route via tailscale0. |
765 | FORWARD RELATED, ESTABLISHED missing | **Fixed** | Installed 'iptables' via apk in 'docker-compose.yml'.
    Rule now injects successfully. |
766 | DOCKER_GW variable parsing error | **Fixed** | Added 'head -1' in 'scripts/routing-fix.sh'. |
767 | 'FIREWALL_OUTBOUND_SUBNETS' was the root cause | **Fixed** | Removed from 'docker-compose.yml'. Gluetun
    no longer hijacks CGNAT routing. |
768 | Host tailscaled error state from aggressive 'tailscale down' | **Mitigated** | Toggle script now
    detects and auto-restarts. |
769
770 ##### 15.1.2 Container Default Route Bypasses WARP
771 **Problem:** Inside the WARP container's network namespace, the main routing table's default route was
    via 'eth0' (Docker bridge), not 'tun0' (WARP tunnel). While gluetun uses policy routing (rule 101
    table 51820 'tun0') for most traffic, traffic originating from the container's own IP
    ('172.18.0.2') hits rule 100 ('from 172.18.0.2 lookup 200'), whose default was also via 'eth0'. This
    caused the SOCKS5 proxy's outbound connections to bypass WARP.
772
773 **Fix:** Updated the 'routing-fix' container to:
774 1. Change the main table's default route to 'default dev tun0'
775 2. Change table 200's default route to 'default dev tun0'
776 3. Add a specific route for the WARP WireGuard endpoint ('162.159.192.1') through 'eth0' via the Docker
    gateway in table 200, preventing a tunnel loop
777 4. Re-assert all routes every 30 seconds (gluetun may reset them on health checks)
778 5. Dynamically detect the Docker subnet from 'ip route show dev eth0' instead of hardcoding
    '172.18.0.0/16'
779
780 ##### 15.1.3 Hardcoded Docker Subnet in Routing Fix
781 **Problem:** The 'routing-fix' container hardcoded '172.18.0.0/16' and '172.18.0.1' for the Docker bridge
    subnet and gateway. If Docker's bridge network changed (after 'docker network prune', Colima
    restart, or on a different machine), these routes would be wrong.
782
783 **Fix:** Detection is now dynamic using 'ip route show':
784 - 'DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')"
785 - 'DOCKER_GW=$(ip route show default | awk '{print $3}')"
786
787 This extracts the actual subnet and gateway directly from the routing table, working with any subnet size
    ('/16', '/24', '/20', etc.).
788
789 ### 15.2 Routing Loops & Subnet Collisions
790
791 ##### 15.2.1 The Infinite Recursive Routing Loop (Hotspot Paradox)
792 **Problem:** A critical network topology crash occurs if the host Mac connects to a Mobile Hotspot (e.g.,
    a Windows PC or Phone) that is *also* connected to the Tailscale mesh and has "Use Exit Node"
    enabled. Because the Mac is hosting the Exit Node, it routes its internet traffic to the Hotspot.
    The Hotspot intercepts the Mac's traffic on its way to the cellular tower and routes it *back* to
    the Exit Node (the Mac) via Tailscale. The Mac receives it, tries to send it out again, and the
    Hotspot intercepts it again. This creates an infinite, recursive encryption loop that instantly
    destroys network bandwidth and drops all connections.
793
794 **Fix:** This is a physical routing paradox that cannot be resolved in software. The upstream physical
    router providing internet to the Mac *must* be allowed to route its traffic directly to the
    physical WAN interface. The user must manually disable "Use Exit Node" on the upstream device
    providing the hotspot.
795
796 **Advanced Workaround:** If the upstream hotspot is a Windows machine and the user explicitly wants it to
    remain connected to the Exit Node, a permanent static route can be injected into the Windows
    routing kernel to forcefully bypass the Tailscale WFP interception for WARP packets. By binding the
    route to the physical adapter's Interface Index (IF), it becomes a permanent, roaming-aware bypass.
797
798 On Windows (Admin Command Prompt):
799 '''cmd
800 route print (find Interface List, note the IF number for the active Wi-Fi adapter, e.g. 15)
801 route -p add 162.159.192.1 mask 255.255.255.255 0.0.0.0 IF 15
802 route -p add 162.159.193.1 mask 255.255.255.255 0.0.0.0 IF 15
803 '''
804
805 On Linux (Root Terminal):
806 '''bash
807 ip route add 162.159.192.1 dev wlan0

```



```

808 ip route add 162.159.193.1 dev wlan0
809 '''
810
811 On macOS (Root Terminal):
812 '''bash
813 route add -host 162.159.192.1 -interface en0
814 route add -host 162.159.193.1 -interface en0
815 '''
816
817 *(Note: If the upstream router is an iPhone or Android device, you cannot inject static routes without
818 jailbreak/root. You must simply disable "Use Exit Node" in the mobile Tailscale app).*
819
820 This punches a tiny hole straight through the paradox, letting the Mac's WARP packets escape to the
821 physical internet while keeping the rest of the upstream router's traffic securely trapped in the
822 Tailscale Exit Node.
823
824 ##### 15.2.2 Infinite Egress Routing Loops & Subnet Takeover Paradoxes
825 **Problem:** Two distinct network loops can break host egress connectivity entirely when the exit node is
826 active:
827
828 1. **The Recursive Host-VM Tunnel Loop:**
829 - **Mechanism:** The host Mac's default route points to the Tailscale exit node ('utun*'). The exit
830 node container sends its encrypted tunnel packets (destined for the WARP endpoint '162.159.192.1')
831 back to the host via the VM NAT. Without a static route exception on the host Mac, the host Mac
832 routes those packets right back into the exit node ('utun*'), creating an infinite loop that
833 blackouts all host network egress.
834 - **ARP-Scoping Pitfall:** Simply binding the bypass route to the interface ('route add -host
835 162.159.192.1 -interface en0') fails when the default route is deleted or redirected to 'utun*'.
836 Since '162.159.192.1' is not local to the physical link, macOS fails to resolve it via ARP, dropping
837 all tunnel packets.
838 - **Fix:** Dynamically resolve the active physical gateway IP (e.g. '172.17.0.1') on startup, and add
839 the static host routes for the WARP endpoints via the gateway IP directly ('route add -host
840 162.159.192.1 <gateway_ip>'). Additionally, because standalone 'tailscaled' on macOS does not
841 automatically modify the system default route via the CLI, we manually redirect the default gateway
842 to 'utun*' using 'setup_exit_node_routing'.
843
844 2. **The Docker Compose Subnet Takeover Collision:**
845 - **Mechanism:** To avoid collisions with the host's campus network, Colima's default bridge ('docker.
846 bip') is moved (e.g. to '10.200.0.1/24'). However, because '172.17.0.0/16' is now free inside the VM
847 , Docker Compose's IPAM dynamically takes it over for the project's custom network ('
848 nullexit_default'). Because the containers receive IPs on '172.17.0.0/16', the host Mac cannot send
849 local peer discovery packets (e.g., Tailscale 'magicsock' packets on '172.17.0.2:41641') directly.
850 The host routing table sends them out 'en0' to the physical Wi-Fi, returning 'host is down' or 'no
851 route to host'.
852 - **Fix:** Explicitly configure a custom default network subnet '10.200.1.0/24' in 'docker-compose.yml
853 ' to prevent Docker Compose from ever reclaiming the conflicting '172.17/172.18' ranges.
854
855 > See also 15.2.1 (The Hotspot Paradox) for the related infinite loop that occurs when the physical
856 upstream router is also a Tailscale exit-node client.
857
858 ##### 15.2.3 Docker Default Bridge vs Host Wi-Fi Subnet Collision (The 172.17.0.0/16 Subnet Overlap)
859 **Context:** When attempting to ping a local Windows client ('172.17.22.187') or the default gateway
860 ('172.17.0.1') directly over Wi-Fi, the Mac host returns 'No route to host' (displaying a 'REJECT' /
861 '!' flag in the routing table). Tailscale on the client also drops offline shortly after starting
862 when using the exit node, failing to establish direct P2P connections and falling back to slow DERP
863 relays.
864
865 **Root Cause & Subnet Overlap Mechanism:**
866 1. **Docker Default Subnet:** By default, the Docker daemon initializes its default bridge network ('
867 docker0') on the **172.17.0.0/16** IP range.
868 2. **Subnet Conflict:** The host's physical Wi-Fi network happens to be assigned the exact same
869 **172.17.0.0/16** range by the local network router.
870 3. **Routing Clash:** Colima launches a Linux VM on the Mac host to run the Docker engine. Because Docker
871 inside that VM creates 'docker0' and claims the '172.17.0.0/16' subnet, any packet sent to '172.17.
872 x.x' from within the containers (such as Tailscale local discovery) is captured by the VM's internal
873 virtual bridge interface (which is down/empty) and is blackholed. On the Mac host, macOS
874 dynamically marks routes as 'REJECT' ('!') when local ARP resolution fails, causing local pings to
875 immediately abort with 'No route to host'.
876 4. **Tailscale Relay Saturation:** Because local peer-to-peer discovery is blackholed, Tailscale falls
877 back to public DERP relay servers (e.g. 'relay "tor"' in Toronto). When exit-node routing is enabled
878 , all client web traffic is routed through the relay. Public relays impose strict rate limits and
879 throttle high-volume traffic, resulting in packet drops. This saturation drops Tailscale's control
880 packets, causing the connection to time out and showing the client as offline.
881
882 **Fix:**
883 The canonical script for this is 'scripts/fix-docker-bridge-collision.sh'. **One command applies the
884 entire fix:**
885
886 '''bash
887 bash scripts/fix-docker-bridge-collision.sh # apply
888 bash scripts/fix-docker-bridge-collision.sh --dry # preview changes without applying

```



```

850 '''
851
852 It auto-detects the host's actual Wi-Fi subnet (no more guessing whether the campus DHCP handed out '172.
      x', '10.x', or '192.168.x'), picks a non-conflicting 'docker.bip' from a small candidate set (
      defaulting to '172.26.0.1/24' if no overlap is matched), backs up '~/.colima/default/colima.yaml'
      with an ISO-8601 timestamp, edits the file **in place** (idempotent replaces an existing 'docker.
      bip' if present, appends under an existing 'docker:' block, or creates a fresh block), restarts
      Colima, rebinds Wi-Fi DHCP, verifies the new default route is no longer Docker's bridge, and re-runs
      both './toggle.sh' and 'scripts/diagnose-host-leak.sh' to print the post-fix verdict.
853
854 For reference, the underlying manual fix is: edit the Colima configuration file ('~/.colima/default/
      colima.yaml') on your Mac to define:
855 '''yaml
856 docker:
857     bip: 172.26.0.1/24    # any /24 that does NOT overlap your Wi-Fi subnet
858 '''
859 Then, restart Colima to apply the new subnet inside the VM:
860 '''bash
861 colima restart
862 '''
863 This forces Docker to use '172.26.x.x' for its default bridge, freeing the physical '172.17.x.x' range
      and letting pings and direct Tailscale peer-to-peer connections flow normally.
864
865 ##### 15.2.4 July 4, 2026: The "Network Unreachable" Host Leak & Post-Wake Container Destructions
866 **Symptom:**
867 1. Running 'docker compose config' with dummy keys corrupted the user's '.env' file by appending invalid
      WireGuard keys. This caused the 'warp' container to become unhealthy on boot, which triggered a full
      teardown of the gateway by 'toggle.sh'.
868 2. After fixing '.env', 'toggle.sh' successfully booted the gateway, but the host's internet started
      bypassing the tunnel entirely (a massive host leak). The Cloudflare trace returned 'warp=off' and
      exposed the host's physical ISP IP.
869 3. Running 'recover.sh --post-wake' would violently force-recreate all gateway containers and drop the
      host's internet connection.
870
871 **Root Cause:**
872 - **Bug 1 (The Host Leak):** In 'toggle.sh', the logic to find the Tailscale 'utun' interface used '
      ifconfig | grep -B4 "inet 100."'. Due to the 4 lines of before-context, this regex was accidentally
      capturing the interface *above* the actual Tailscale interface (e.g. grabbing 'utun4' instead of '
      utun5'). Because 'utun4' had no IP address, the subsequent 'sudo route add -net 0.0.0.0/1 -interface
      utun4' silently failed with "Network is unreachable", leaving the host's default route exposed.
873 - **Bug 2 (The Post-Wake Destructions):** The health check in 'recover.sh --post-wake' relied on 'docker
      compose exec -T warp curl -sf https://www.cloudflare.com/cdn-cgi/trace'. However, the newer 'qmcgaw/
      gluetun' Alpine-based Docker image no longer ships with 'curl' pre-installed. The health check
      failed with 'executable file not found in $PATH', tricking 'recover.sh' into thinking the tunnel was
      completely dead and needed to be forcefully recreated via 'docker compose up -d --force-recreate'.
874 - **Bug 3 (Restart Flags):** Previous refactors incorrectly permitted 'toggle.sh restart' instead of
      strictly enforcing 'toggle.sh --restart'.
875
876 **Resolution:**
877 - Cleaned the '.env' file of duplicate dummy entries.
878 - Replaced the brittle 'grep -B4' interface parsing in 'toggle.sh' with a strict AWK parser: 'awk '/^[a-
      z0-9]+/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':''. This mathematically isolates the
      exact interface possessing the Tailscale '100.x.x.x' IP address, preventing the '0.0.0.0/1' route
      from silently failing.
879 - Changed the health check in 'recover.sh' to use 'wget -q0- --timeout=3' which is natively installed in
      the Gluetun image. This stopped the unnecessary destructive recreations of the containers during
      post-wake events.
880 - Reverted the 'restart' alias in 'toggle.sh' to strictly require '--restart'.
881
882 ##### 15.2.5 July 4-5, 2026: Tailscale Data-Plane Loop (The DERP Relay Deadlock)
883 **Symptom:**
884 After bypassing the Tailscale control plane ('192.200.0.0/16') to fix the "offline" mesh status, the Mac
      appeared online. However, external peer devices (like the user's phone on cellular) could not
      successfully ping or SSH into the Mac. 'tailscale netcheck' reported 'Nearest DERP: unknown (no
      response to latency probes)'.
885
886 **Root Cause:**
887 While bypassing the control plane kept the daemon connected to the coordination servers, Tailscale's
      actual peer-to-peer data plane relies on STUN (for NAT traversal) and DERP relays. These relays span
      ~80 dynamically allocated IP addresses across multiple cloud providers. Because we were forcing
      '0.0.0.0/1' through the 'utun*' exit node, the Mac's own outbound connection attempts to DERP relays
      were being routed back into the tunnel. To prevent an infinite loop, Tailscale's daemon dropped its
      own packets when it saw them returning on the TUN interface. This effectively left the daemon deaf
      and blind to peer connections.
888
889 **Resolution:**
890 - Modified 'add_warp_bypass_routes' in 'scripts/common.sh' to execute a live API fetch ('curl -s https://
      login.tailscale.com/derpmap/default') during startup.
891 - The script parses out all active DERP relay IPv4 addresses (~80 IPs) and adds a physical host bypass
      route for every single one of them.

```

```

892 - These IPs are written to a temporary file ('/tmp/nullexit-derp-ips.txt') so they can be cleanly un-
      routed by 'remove_warp_bypass_routes' when the gateway shuts down. Peer connectivity (ping, SSH,
      SFTP) was fully restored.
893
894 ##### 15.2.6 July 5, 2026: Hardcoded WARP Endpoints in docker-compose
895 **Symptom:**
896 The bash scripts allowed overriding the default Cloudflare WARP IP endpoints via '.env' variables ('
      WARP_ENDPOINT_1' and 'WARP_ENDPOINT_2') to establish bypass routes. However, 'docker-compose.yml'
      statically hardcoded '162.159.192.1' for the 'warp' container's 'WIREGUARD_ENDPOINT_IP'.
897
898 **Root Cause & Risk:**
899 If a user set a custom endpoint in '.env', the script would successfully bypass the custom IP on the host
      level. However, Gluetun would still attempt to connect to the hardcoded '162.159.192.1'. Since
      '162.159.192.1' was no longer explicitly bypassed, its packets would be sucked into the 'utun*' exit
      node, causing an infinite WireGuard routing loop that instantly breaks the gateway.
900
901 **Resolution:**
902 Modified 'docker-compose.yml' to dynamically read the environment variable via '${WARP_ENDPOINT_1
      :-162.159.192.1}', ensuring both the host routing scripts and the container use the exact same
      endpoint.
903
904 ##### 15.2.7 Incident Post-Mortem: WARP Watcher Race vs Post-Wake (July 10, 2026)
905 **Symptom:** After switching Wi-Fi networks, the host IP appeared as the raw ISP IP ('132.x.x.x') instead
      of a Cloudflare/WARP IP. The gateway appeared to complete post-wake recovery successfully in the
      logs, but nuclear shutdown fired immediately after and killed everything.
906
907 **Root Cause:** A fatal race condition between two concurrent processes:
908
909 1. **Wi-Fi roam** 'watcher.sh' fires 'recover.sh --post-wake'
910 2. **Post-wake** finds the 'warp' container missing/dead calls 'docker compose up -d --force-recreate'
      (~35 seconds)
911 3. **In parallel**, the WARP Watcher (running as a child of 'toggle.sh') polls the warp container every 5
      s during failures
912 4. While the container is being recreated it returns 'warp=off' for those 35 seconds 7 consecutive
      failures **WARP SHUTDOWN fires**, calling nuclear 'recover.sh'
913 5. Nuclear 'recover.sh' tears down Tailscale, resets DNS to 1.1.1.1, stops all containers **gateway dead
      , IP exposed**
914 6. The post-wake 'docker compose up' finishes at ~35s mark, container healthy but the gateway is already
      gone
915
916 The evidence in 'output.log':
917 ```
918 [2026-07-10T06:07:27Z] NET: ... bash recover.sh --post-wake
919 [2026-07-10T06:07:47Z] WARP DOWN cdn-cgi/trace reports warp=off watcher starts counting
920 ...
921 add net 0.0.0.0: gateway utun5 post-wake successfully re-routes at 02:08:22
922 ...
923 [2026-07-10T06:09:04Z] WARP SHUTDOWN 6 consecutive failures watcher fires nuclear
924 ```
925
926 **Fix (July 10, 2026):** Added a **WARP Watcher inhibit marker** ('/tmp/nullexit-warp-inhibit.marker'):
927 - 'recover.sh --post-wake' writes the marker at startup **before** any container operations
928 - The WARP Watcher loop checks for the marker at the top of every iteration; if present it resets '
      consec_off=0' and sleeps 5s (skips the poll entirely)
929 - 'recover.sh' removes the marker at the very end (on both success and failure via 'trap cleanup_inhibit
      EXIT')
930 - This guarantees the watcher never counts failures during the intentional container-down window of a
      post-wake force-recreate
931
932 The marker file path is hardcoded to '/tmp/nullexit-warp-inhibit.marker' in both 'toggle.sh' and 'recover
      .sh'.
933
934 ##### 15.2.8 Incident Post-Mortem: Concurrency Collision between Nuclear Teardown and Post-Wake (July 10,
      2026)
935 **Symptom:** When the background WARP Watcher triggers a nuclear shutdown (due to a real or simulated
      tunnel failure), the gateway is completely torn down. However, the network configuration changes
      made during the teardown trigger a 'State:/Network/Global/IPv4' network change event. The
      LaunchAgent-based network watcher ('watcher.sh') intercepts this event and concurrently spawns '
      recover.sh --post-wake' to heal/restore the gateway. This causes 'docker compose down' (from nuclear
      teardown) and 'docker compose up' (from post-wake) to run in parallel, resulting in container
      errors like 'dependency failed to start: container warp exited (0)' and the gateway becoming wedged.
936
937 **Root Cause:** A lack of coordination/locking between the lifecycle control scripts. While 'toggle.sh'
      used '/tmp/nullexit-toggle.lock' to prevent concurrent 'toggle.sh' runs, 'recover.sh' (both nuclear
      and '--post-wake' modes) did not utilize or check any lock files.
938
939 **Fix (July 10, 2026):** Implemented a shared lifecycle lock file ('/tmp/nullexit-toggle.lock') across
      both scripts:
940 - **'recover.sh' (all modes):** Checks '/tmp/nullexit-toggle.lock' at startup. If the lock is held by
      another running lifecycle script ('toggle.sh' or 'recover.sh'):

```

941 - In '--post-wake' mode, it logs a skip message to 'output.log' and exits cleanly ('exit 0'). This
safely prevents the auto-recovery agent from recreating containers during a teardown.

942 - In nuclear recovery mode, it exits with an error.

943 - *****'recover.sh' Lock Cleanup**:** Cleans up the lock file upon completion using 'trap cleanup_recover EXIT'
(which checks if the lock file belongs to the current PID).

944 - *****'toggle.sh' update**:** The lock check in 'toggle.sh' was updated to look for both 'toggle.sh' and '
recover.sh' in the running processes to enforce proper exclusion.

945

946 **#### 15.2.9 Incident Post-Mortem: Infinite Auto-Recovery Loop after WARP Watcher Nuclear Shutdown (July
10, 2026)**

947 ****Symptom**:** When the background WARP Watcher triggered a nuclear shutdown (due to a tunnel outage), it
successfully executed 'recover.sh' (nuclear). However, immediately after the teardown completed, the
system started spawning 'recover.sh --post-wake' and rebuilding the containers again. This created
an infinite loop of tearing down and auto-restarting the gateway, keeping the host's internet route
permanently wedged.

948

949 ****Root Cause**:** An active-state marker file leak. The LaunchAgent-based network watcher ('watcher.sh')
only fires 'recover.sh --post-wake' when '/tmp/nullexit-gateway-active.marker' is present on disk.
While 'toggle.sh' (STOP path) correctly deleted this marker, the nuclear 'recover.sh' script did not
. When the WARP Watcher ran 'recover.sh' (nuclear), the containers were stopped but the active-state
marker was left on disk. The network changes from the teardown then triggered 'watcher.sh', which
saw the marker, believed the gateway was still supposed to be active, and triggered '--post-wake' to
start it back up.

950

951 ****Fix (July 10, 2026)**:** Modified the start of 'recover.sh' to explicitly delete '/tmp/nullexit-gateway-
active.marker' if 'POST_WAKE' is 'false' (nuclear mode). This ensures that any subsequent network
configuration changes made during the teardown are correctly ignored by 'watcher.sh' because the
marker file is gone.

952

953 **### 15.3 MTU, Fragmentation & Throughput**

954

955 **#### 15.3.1 Network Bufferbloat & MTU Optimizations (Double-VPN UDP Crash)**

956 ****Symptom**:** When running aggressive UDP bandwidth tests (such as macOS 'networkQuality', which uses QUIC
/HTTP3), the 'nullexit' tunnel completely crashes or exhibits severe lag (6,000ms+ bufferbloat).

957

958 ****Root Cause**:** The bug is a cascading MTU mismatch. Tailscale generates a UDP packet. Cloudflare WARP (
also WireGuard) receives the packet, wraps it in another layer of encryption, pushing the packet
size beyond the standard 1500 byte limit. Unlike TCP, UDP packets cannot be resized dynamically via
MSS negotiation, resulting in massive IP packet fragmentation. The Apple Virtualization framework ('
vz'), which Colima uses for bridged networking, silently drops heavily fragmented UDP packets under
high load. This blackholes the entire UDP upload stream, instantly dropping the connection.

959

960 ****Fix (Partial)**:** To prevent MTU fragmentation for standard web traffic, ****TCP MSS Clamping**** is
strictly enforced via the macOS 'pf' firewall. By adding 'scrub out all max-mss 1160' to 'scripts/pf
.conf', macOS unilaterally shrinks all outgoing TCP headers to fit comfortably inside the double-
encrypted tunnel, bypassing Path MTU Discovery blackholes. 'TS_DEBUG_MTU=1200' was also added to '
docker-compose.yml' to minimize internal fragmentation. *(Superseded see 15.3.2: 1160 was later
found still too large for double-tunneled traffic and tightened to 1120. That fix originally only
reached the container's 'post-rules.txt'; 'pf.conf's host-side clamp stayed stale at 1160 until it
was caught and synced to read 'GATEWAY_MSS' dynamically, same as 'post-rules.txt', so the host and
container can no longer drift apart.)*

961 ****Unresolved**:** The only way to solve the UDP crash bug natively is to artificially cap the maximum
tunnel bandwidth using SQM ('fq_codel'). Because a static bandwidth cap would artificially throttle
high-speed networks, 'nullexit' leaves the bandwidth uncapped, optimizing perfectly for TCP while
accepting UDP fragmentation under extreme edge-case load tests.

962 ****Why not Dynamic SQM (Aurate)?**** Implementing an EMA/PID-controlled aurate daemon (like 'sqm-aurate'
) requires a constant 100ms polling loop to measure the kernel packet queue, which severely
degrades laptop battery life by preventing deep C-state CPU idling. Furthermore, macOS's native
firewall ('dummynet') only supports 'fq_codel' and does not support the newer Linux 'CAKE' algorithm
, which is the only algorithm that offers kernel-native, event-driven 'aurate-ingress' (zero
polling penalty). Therefore, dynamic SQM is computationally hostile to battery-powered macOS hosts.

963

964 **#### 15.3.2 Double-Tunneling MSS Clamping (Stalling Bug)**

965 ****Problem**:** Traffic double-wrapped through Tailscale (WireGuard) and WARP (WireGuard) suffered 120 bytes
of MTU overhead (60 + 60). The default MSS clamp of 1180 was still slightly too large, causing
large packets to fragment or drop, which resulted in mysterious web page stalls on strict-MSS
endpoints.

966

967 ****Fix**:** Lowered the TCP MSS clamp in 'post-rules.txt' to '1120' to guarantee double-tunneled payloads
fit safely within standard internet MTUs.

968

969 **#### 15.3.3 Wi-Fi Edge Packet Loss via QUIC (UDP) Fragmentation**

970 ****Problem**:** Applications that heavily utilize HTTP/3 (QUIC) over UDP such as Facebook, Instagram, and
Google services experience massive stuttering and stalled image loading when the client device is far
away from the router (weak Wi-Fi signal). The issue did not occur when close to the router.

971

972 ****Root Cause**:** The 'nullexit' gateway uses a double-encrypted tunnel (Tailscale + WARP). To prevent
packets from fragmenting when entering these tunnels, a firewall rule in 'post-rules.txt' uses '--
set-mss 1120' to safely clamp TCP packets. However, because QUIC runs entirely over UDP, it bypasses
the TCP MSS handshake completely. QUIC blasts maximum-MTU (1280 byte) UDP packets. The inner '
tailscale0' interface (MTU 1200) forces these large UDP packets to fragment into two pieces. When

the client is on the edge of Wi-Fi range (where 5-10% packet loss is common), losing *either* UDP fragment destroys the *entire* image frame. This exponential amplification of packet loss stalled QUIC streams.

****Fix & Trade-offs:**** Appended an 'iptables' rule to 'post-rules.txt' that explicitly blocks 'UDP --dport 443' (QUIC). When modern web apps detect QUIC is blocked, they instantly fall back to HTTP/2 (TCP 443). Because TCP is routed through the MSS clamp, the packets are resized *before* they leave, entirely eliminating IP fragmentation and restoring high-speed loading over weak Wi-Fi.

Consequences: This adds a minor latency penalty (TCP 3-way handshake) compared to QUIC's zero-RTT connection. It may also force WebRTC video calls (Google Meet/Discord) to fall back to TCP TURN servers, slightly inflating real-time video latency. However, in a constrained double-tunnel mesh, the absolute stability gained from eliminating fragmentation vastly outweighs these trade-offs.

15.3.4 Cellular Networks & PMTUD Blackholes (The 1280 Byte Limit)

****Experiment:**** Conducted a live packet sweep from the PC-1 host to Phone-1 connected via a cellular network + Tailscale DERP relay using 'ping -D -s <size>'.

****Result:****

- Packets with a payload of '1252' bytes (+ 28 bytes IP/ICMP headers = ****1280 bytes total****) succeeded perfectly with ~60ms latency.
- Packets with a payload of '1253' bytes (****1281 bytes total****) and above resulted in a silent timeout.

****Analysis:**** The cellular network (or the Tailscale DERP relay) enforces a strict MTU of 1280 bytes. Critically, it acts as a "PMTUD Blackhole" it silently drops packets larger than 1280 bytes instead of returning an ICMP "Fragmentation Needed" warning. If the exit node tries to send a standard 1500-byte internet packet to the phone over this link, it vanishes, causing web pages to stall permanently.

****Conclusion:**** This empirically validates the absolute necessity of the TCP MSS clamping rule in 'post-rules.txt'. By artificially clamping the MSS to '1120' (the value 15.3.2 above found necessary '1180' measured too large), we ensure the TCP payload + headers + double-WireGuard overhead never exceeds the strict 1280-byte ceiling of the cellular mesh link.

15.3.5 Low Gateway Throughput (DERP Relay & MTU Fragmentation) Fixed

****Problem:**** The host Mac's internet throughput through the gateway dropped significantly (e.g., from 34 Mbps raw to 2.1Mbps via gateway). This was caused by a combination of two bottlenecks:

****1. Tailscale DERP Relay Bottleneck:****

Because the host Mac and the Docker container operate behind the same public IP, standard hole-punching for Tailscale P2P failed, forcing traffic through Tailscale's NYC DERP relay. Normally, 'routing-fix.sh' explicitly drops container Tailscale UDP packets destined for local subnets (when 'TAILSCALE_ALLOW_LAN_P2P=false' in '.env' or auto-detected). We nuke these P2P connections on purpose to prevent the severe macOS 'gvproxy' SNAT endpoint poisoning issue we faced earlier (see 15.7 SNAT Endpoint Poisoning). However, this strict policy unintentionally blocked direct communication between the container and the Mac host itself via the Docker bridge network.

****Fix:**** Tailscale's control plane registers the Mac host using its physical IP (e.g., '172.17.52.223'), not the Docker bridge IP. If this physical IP falls within the dropped local subnets (like '172.16.0.0/12'), the P2P handshake is dropped. To solve this natively, 'scripts/common.sh' now features a 'write_host_ips()' function that actively grabs all physical IPv4 addresses on the Mac and writes them to '.host_ips'. 'routing-fix.sh' dynamically reads this file every 30 seconds and injects priority 'RETURN' rules for those exact physical IPs. This guarantees direct container-to-host P2P over the local bridge (bypassing DERP) while continuing to block external LAN devices (like phones) to prevent the 15.7 SNAT poisoning bug.

****2. MTU Double-Encapsulation Fragmentation:****

The host's Tailscale interface ('utun5') defaulted to an MTU of 1280. The WARP tunnel inside the container also uses an MTU of 1280. When a full 1280-byte Tailscale packet from the host entered the container and was encapsulated by WARP (adding ~60 bytes of overhead), the resulting packet exceeded 1280 bytes, leading to severe fragmentation and performance degradation.

****Fix:**** In 'toggle.sh's 'setup_exit_node_routing' function, the host's Tailscale 'utun' interface MTU is explicitly lowered to '1200' (configurable via 'HOST_MTU' in '.env'). This ensures that host-originated packets can comfortably fit inside the container's WARP tunnel encapsulation without fragmenting.

> **Design Note Is direct hostcontainer P2P a Docker bug, or are we "exploiting a VM/host relationship"?

> Neither. It is expected, well-defined behavior, and the RETURN-rule carve-out is a *defensive de-restriction*, not an exploit. The reasoning:

- ****On native Linux****, the container and host share one kernel; the Docker bridge is a first-class local interface and hostcontainer is **trivially** direct. Nobody calls that an exploit. Our RETURN rules simply reproduce that same local reachability on macOS.
- ****The RETURN rules add nothing new; they *stop blocking* something legitimate.**** 'routing-fix.sh' broadly DROPS containerLAN Tailscale UDP to prevent the real 15.7 'gvproxy' SNAT endpoint-poisoning bug caused by **external** LAN peers (e.g. a phone on another AP). Tailscale's control plane registers the Mac host by its **physical** LAN IP (e.g. '172.17.52.223'), which falls inside those DROPPed ranges so the host, the one peer that is literally the same physical machine, became collateral damage. The '.host_ips' RETURN rules are the minimal, precise exception that re-permits exactly that same-machine path and nothing else.

```

1005 > - **The genuinely VM-specific wart is 'gvproxy' SNAT poisoning (15.7), and we route *around* it, not
      through a hole in it.** Keeping hostcontainer traffic on the Docker bridge (direct) avoids letting
      it get NATed through the userspace 'gvproxy' endpoint rewriter that confuses Tailscale's peer-
1006 endpoint learning. That is the opposite of exploiting the VM boundary it is respecting it.
      > - **What genuinely *doesn't* work on a VM host is UDP hole-punching to arbitrary *external* peers** (
        see 15.13). That is a real limitation of userspace VM networking, and it is why remote peers fall
        back to DERP unless bridged mode is enabled on a trusted LAN. Direct hostcontainer P2P is the one
        case that works regardless because both endpoints live on the same physical machine, no hole-
        punching is needed at all. It was verified 'direct 172.17.52.223' even in plain user-mode networking
        (no '--network-address').
1007 >
1008 > In short: the DROP is the deliberate policy, the RETURN is a surgical exception for same-machine
      traffic, and the whole thing is standard bridge networking not a Docker bug and not an exploit.
1009
1010 ##### 15.3.6 The "~4Mbps Ceiling" Was a Benchmark Artifact, Not a Tunnel Bottleneck (July 18, 2026)
1011 **Symptom:** Repeated throughput tests against 'speedtest.tele2.net' from both the host (full double-
      tunnel) and the 'warp' container (isolated via its own SOCKS5 proxy) consistently landed around
      3.3-4.7 Mbps, regardless of which leg was tested. This looked like a hard nullexit-imposed ceiling
      and was the leading suspect for the Instagram/general-slowness symptoms in 15.3.3-15.3.4.
1012
1013 **Investigation:** Before accepting "the tunnel caps at ~4Mbps," the same target was measured **raw**
      gateway fully stopped, plain ISP egress, no tunnel at all and it *still* came back at ~3.2 Mbps. A
      second raw test against a fast, nearby CDN (Cloudflare) returned ~32.2 Mbps on the same physical
      connection. 'speedtest.tele2.net' was the bottleneck, not nullexit; it's a slow/distant target that
      happened to be the first thing tried.
1014
1015 **Confirmation:** Re-measured through the tunnel against a fast CDN that (unlike Cloudflare's own speed-
      test endpoint, which 403s WARP's shared/CGNAT egress IPs as abuse-prevention) doesn't flag WARP IPs
      Hetzner's public speed-test files. Result: **15.8-26.4 Mbps** tunneled, in the same ballpark as the
      32.2 Mbps raw baseline (roughly 20-50% overhead from the double WireGuard encapsulation, normal for
      this architecture). CPU and memory on the Colima VM were also checked during sustained transfers ('
      colima ssh -- top -bn1' / 'free -m') and ruled out as a bottleneck load average stayed near
      0.00-0.07 and swap held flat around 250MB throughout, on both a run that completed cleanly and one
      that reset mid-transfer.
1016
1017 **Fix:** 'scripts/sweep.sh's throughput check (9, check 7/7) now defaults to the Hetzner target instead
      of 'speedtest.tele2.net', so future sweeps report real tunnel performance instead of accidentally re-
      measuring a slow third-party server and misattributing it to nullexit.
1018
1019 **Residual, unchanged:** the occasional mid-transfer connection reset (curl exit 56, 'Recv failure:
      Connection reset by peer') observed on maybe 1-in-6 to 1-in-8 attempts is still unexplained by CPU/
      memory and remains consistent with the gvproxy/UDP-fragmentation-under-load issue in 15.3.1 a real,
      separate, and still-open question from the throughput-ceiling one this section resolves.
1020
1021 ### 15.4 DNS Interception & Proxy
1022
1023 ##### 15.4.1 DNS Proxy: TCP Wire Format Mismatch (socat / dnsmasq)
1024 **Problem:** When the Tailscale data plane is unavailable (DERP relay only), the script falls back to a
      local DNS proxy that forwards queries from host UDP:53 to Docker's TCP:5354 (AdGuard). The 'socat'
      and 'dnsmasq' approaches both failed.
1025
1026 - **socat** forwards raw bytes it does not prepend the 2-byte length prefix that DNS-over-TCP requires.
      AdGuard parses the stream incorrectly and returns garbage.
1027 - **dnsmasq** tries UDP first to reach the upstream ('127.0.0.1#5354'), but Colima's SSH tunnel only
      forwards TCP. With a single upstream server, dnsmasq doesn't fall back to TCP even with '--timeout
      =1'.
1028
1029 **Solution:** Replaced both with a **Python DNS proxy** ('scripts/dns-proxy.py', ~25 lines) that properly
      handles the DNS-over-TCP wire format: reads a UDP query from the host, prepends the 2-byte length
      prefix, sends it over TCP to AdGuard, strips the prefix from the response, and sends it back over
      UDP.
1030
1031 *Architectural Note: Why is the DNS proxy in Python while the SOCKS5 proxy was moved to a compiled Go
      Docker container?*
1032 The SOCKS5 proxy handles heavy data streaming (e.g., video, downloads). Python introduces massive CPU/
      battery overhead for raw socket streaming, which is why we rely on the compiled Go implementation
      built directly into the 'tailscaled' daemon via 'TS_SOCKS5_SERVER=:1080'.
1033 Conversely, the DNS proxy only handles intermittent UDP queries (a few bytes per minute) generated solely
      by the local host machine. The Python overhead for this is effectively 0.001 Watts. We deliberately
      kept 'dns-proxy.py' in Python because it runs natively on the macOS/Linux host rewriting it in Go
      would force users to install a Go compiler ('brew install go') on their host machine for zero
      noticeable battery gain. (Note: If this architecture is ever adapted to serve as a public DNS
      resolver for thousands of users or for massive automated web-scraping clusters, 'dns-proxy.py' must
      be rewritten in Go to survive the concurrent packet flood).
1034
1035 ##### 15.4.2 Tailscale MagicDNS Bypassing the AdGuard PREROUTING Intercept
1036 **Context:** The 'routing-fix.sh' script applies an iptables NAT rule ('PREROUTING -i tailscale0 -p udp
      --dport 53 -j REDIRECT') to intercept all incoming DNS queries from connected Tailscale peers and
      force them into the AdGuard container on port 5335.
1037 **Problem:** When remote clients (like Android or iOS) have "Use Tailscale DNS settings" turned on, they
      query Tailscale's MagicDNS IP ('100.100.100.100'). This packet enters the exit node, but the '

```


tailscaled' daemon running *inside* the gateway container intercepts the '100.100.100.100' packet internally. 'tailscaled' then generates a brand new, local DNS request to the upstream DNS server to resolve the query. Because this new request is generated *locally* by the daemon (not arriving externally via the 'tailscaleO' interface), it bypasses the 'PREROUTING' chain completely. The request glides straight out the WARP tunnel to Cloudflare/Google, entirely bypassing the AdGuard sinkhole.

1038 ****Fix:**** The only way to fix this blindspot is to force 'tailscaled' to use AdGuard as its upstream. In the Tailscale Admin Console, the gateway's Tailscale IPv4 address (e.g., '100.x.x.x') must be added as the sole Custom Global Nameserver, and "Override local DNS" must be enabled. Additionally, to block Android devices from bypassing this via Private DNS (DNS-over-TLS), outbound Port 853 traffic is explicitly REJECTED in both 'routing-fix.sh' and 'post-rules.txt'.

1039

1040 **#### 15.4.3 Seamless Wi-Fi Roaming & The macOS DNS Wipeout Loophole**

1041 ****Problem:**** The Tailscale, WireGuard, and Colima NAT stacks are natively designed for connectionless roaming and will seamlessly survive when the macOS host switches Wi-Fi networks (e.g., roaming from home Wi-Fi to a cellular hotspot). However, the internet completely breaks upon network transition. This occurs because macOS intentionally wipes out all custom DNS settings ('networksetup -setdnsservers') and reverts to the new network's default DHCP nameserver whenever the active Wi-Fi BSSID changes. Since the Mac is locked to the exit node, its DNS requests are swallowed by WARP and dumped onto the public internet, which cannot route private DHCP IPs. This results in a total DNS failure.

1042

1043 ****Solution:**** Implementing a silent background "DNS Watcher" daemon ('nullexit-dns-watcher') in 'toggle.sh'. When the gateway starts, it spawns a background polling loop that checks the active Wi-Fi DNS every 30 seconds. If macOS resets the DNS during a network change, the watcher instantly detects the drift and forcefully re-injects '100.100.21.8' via 'networksetup'. When the gateway stops, the watcher process is cleanly killed. This allows true, seamless roaming across physical networks without ever dropping the gateway state.

1044

1045 **#### 15.4.4 docker compose exec '\r' Injection**

1046 ****Fix:**** All 'docker compose exec' output piped through 'tr -d '\r''.

1047

1048 **#### 15.4.5 Stderr Invisibly Swallowed**

1049 ****Symptom:**** Failures were silent due to stderr redirected to 'output.log'.

1050

1051 ****Fix:**** Tailscale wait loop now shows output. Critical commands show errors inline.

1052

1053 **### 15.5 Roaming, Sleep/Wake & Auto-Recovery**

1054

1055 **#### 15.5.1 Gateway Breakage on Lid Close / Wi-Fi Roam Post-Wake + Post-Roam Auto-Recovery**

1056 ****Context / Symptom:**** Closing the lid (sleep/wake) OR losing + reconnecting Wi-Fi (e.g. in an elevator, caf, or roaming between APs) leaves the gateway in a partly-dead state. Symptoms range from a ~30s window of dead DNS to a permanent outage that only full 'recover.sh' or 'toggle.sh' fixes. The root cause is that nullexit has ****no daemon watching macOS sleep/wake or network-state-change events**** the existing DNS Watcher inside 'toggle.sh' only runs in-process and is suspended along with the rest of the user shell on lid close, so the gateway cannot self-heal after either event.

1057

1058 ****Diagnosis:**** Two distinct failure modes share the same root gap.

1059

1060 * **** (A) Lid close sleep wake.**** 'caffeinate -i' (used by 'toggle.sh') blocks *idle* sleep but ****does not block forced sleep from lid close****. Closing the lid hard-suspends Colima VM + every container + every Tailscale-quit-shells-except-tailscaled. On wake the VM clock needs ~5-30s to resync via NTP, during which TLS cert validation fails: Cloudflare WARP '162.159.192.1:2408' rejects the handshake gluetun's healthcheck ('interval: 2s, retries: 15') eventually flips unhealthy Docker 'unless-stopped' restarts the 'warp' container (~30s gap). 'mDNSResponder's in-memory cache still holds pre-sleep entries (stale A records for tailnet hosts) so the first dozen DNS queries after wake time out. 'tailscaled' on the host often keeps its stale DERP relay preference and routes packets through a dead relay for another 30-60s.

1061 * **** (B) Wi-Fi roam / loss in elevators, captive portals.**** WARP is a UDP tunnel to Cloudflare's anycast edge. Carrier NATs and captive-portal networks silently invalidate the UDP binding on roam and on hotspot-bound re-association. macOS 'networksetup' ***should*** preserve DNS settings on the same Wi-Fi service across a roam, but most captive-portal networks DHCP-replace DNS to the captive-portal resolver ****during the captive-portal dance itself****, before the user has authenticated so even DNS hijack survives a clean roam and quietly breaks on captive-portal handoff.

1062

1063 ****Resolution.**** A single ****launchd LaunchAgent**** ('com.nullexit.wake-recovery') runs a long-running shell daemon ('scripts/watcher.sh') that listens for both event sources and, on each fire, executes 'bash recover.sh --post-wake' a ***lighter-touch*** recovery mode that's the opposite of the existing '--nuclear' semantics: it keeps the gateway live while still refreshing every stale subsystem.

1064

1065 ****Deep dive: Apple Power Management Primer (for the unfamiliar).****

1066 macOS exposes several relevant surfaces; the right primitive depends on what you actually need to detect. None of these are obvious and none of them have a standard splash screen in the docs.

1067

1068 * **** 'caffeinate <flags>'**** creates I/O Kit power assertions via 'IOPMAssertionCreateWithName'. It does NOT block forced sleep (lid close, low-battery shutdown, sleep timer firing on a per-set Energy Settings policy). Flags:

1069 * **'-i'** PreventIdleSleep (the only one 'toggle.sh' uses; protects against the OS going idle while the user is active but a clamshell close still hard-puts it to sleep)

1070 * **'-s'** PreventSystemSleep (only honored on AC power; the user may be on battery)

```

1071 * '-u' DeclareUserActive (resets the idle timer every 30s by default; equivalent to keeping the cursor
1072 * '-d' PreventDisplaySleep
1073 * '-m' PreventDiskIdleSleep
1074 * There is NO '-l' (lid-block) flag in any modern macOS. Lid close is a kernel-firmware-level event
1075 * To prove any of this on live hardware: 'pmset -g assertions' lists every active assertion with type /
1076 * process / reason / timeout.
1077 * 'pmset -g log | grep -E 'Sleep|Wake|DarkWake'' shows the recent timeline.
1078 * 'pmset -g history' is the per-day summary.
1079 * 'pmset -g assertions' is the live assertion list (matches '-i -s -d -u' to processes).
1080 * 'launchd' does NOT have a native "on-wake" hook. Not in any version of macOS as of Sonoma. The
1081 * canonical recipe is 'StartCalendarInterval' with a sub-minute cadence and let launchd queue missed
1082 * intervals for replay-on-wake, or use a third-party daemon (Sleepwatcher). For our use case a long-
1083 * running shell daemon listening to the unified log is more responsive than a polling agent.
1084 * 'com.apple.system.powermanagement.* Darwin notifications' ('willSleep', 'didWake', 'didDim', '
1085 * didUndim') are the C-level signal. From a shell, they are best reached through the unified log with
1086 * 'log stream --predicate' see the watcher implementation.
1087 * 'scutil n.watch' is the canonical CLI for live-following network-state changes. Pairs with 'n.add
1088 * State:/Network/Global/IPv4' to get a single tap that fires on **any** IPv4 link/route/DNS/address
1089 * change across **all** interfaces. This is what 'scripts/watcher.sh's network-change listener uses.
1090 * 'com.apple.networkChange' Darwin notification is the C-level equivalent but has no shell-friendly
1091 * listener; use 'scutil' instead.
1092 * 'SCNetworkReachability' is the Objective-C API used by SOCKS5/HTTP clients to detect route changes
1093 * per target. Never a fit for shell scripts.
1094
1095 **The Fix (component-by-component).**
1096
1097 1. 'recover.sh --post-wake' (new flag, **non-destructive** by design). Adding '--post-wake' to the
1098 * existing nuclear recovery script inverts the semantics. The default mode still tears down Tailscale,
1099 * resets DNS to empty, stops caffeinate, runs 'docker compose down', power-cycles Wi-Fi. The new mode
1100 * does:
1101 * Skip Tailscale disconnect (we WANT it to keep the mesh connection)
1102 * **Re-hijack** DNS to the gateway Tailscale IP read from '.gateway_ip' (instead of resetting to empty)
1103 * applies to both 'ACTIVE_SERVICE' and 'ENO_SERVICE' because the system resolver is scoped to en0
1104 * **Refresh** the exit-node preference with 'tailscale up --reset --ssh=true --accept-dns=false --exit-
1105 * node=\"$STS_IP\" --exit-node-allow-lan-access=true' (the exact flags 'toggle.sh' START uses). This re-
1106 * asserts DERP mapping without dropping the mesh
1107 * Inspect 'docker inspect --format '{{.State.Health.Status}}' nullexit-warp-1' and **only** 'docker
1108 * compose up -d --force-recreate warp' if gluetun is unhealthy this single targeted recreate nudges
1109 * the UDP NAT binding back to Cloudflare without disturbing Tailscale or AdGuard
1110 * Run the sharingd-reset step (existing fix from 'toggle.sh' START path; see 15.11) so AirDrop doesn't
1111 * freeze on the new IP lease
1112 * Verify the gateway with 'dig +tcp @127.0.0.1 -p 5354 google.com' (AdGuard via TCP through Colima's
1113 * SSH tunnel) and a 4-second 'tailscale ping' to the gateway Tailscale IP, instead of the default mode
1114 * 's direct-to-1.1.1.1 check
1115 * Crucially: **skip 'sudo -v' because the LaunchAgent has no TTY to prompt on; all privileged calls
1116 * use 'sudo -n' and lean on '/etc/sudoers.d/nullexit' NOPASSWD entries ('networksetup', 'dnsmasq', '
1117 * dscacheutil', 'killall', 'pkill').
1118 2. 'toggle.sh' writes and clears '/tmp/nullexit-gateway-active.marker' (an ISO-8601 UTC timestamp) at
1119 * the bottom of the START path and the top of the STOP path / 'cleanup_handler'. Two new helpers: '
1120 * write_gateway_active_marker' and 'clear_gateway_active_marker'. The watcher uses this file as its **
1121 * only signal** that the gateway is live: if 'toggle.sh' was stopped (or never started), the watcher
1122 * no-ops. This prevents waking-up-only-on-counterfactual events from churning DNS.
1123 3. 'scripts/watcher.sh' (long-running daemon, sibling-style source file). Two backgrounded pipelines:
1124 * **Wake listener**: 'log stream --predicate 'composedMessage CONTAINS[c] \"didWake\" OR composedMessage
1125 * CONTAINS[c] \"Wake from Sleep\" OR composedMessage CONTAINS[c] \"Waking from\" OR composedMessage
1126 * CONTAINS[c] \"DarkWake\"' --style compact --info' piped through 'while read; do run_recover \"WAKE:
1127 * ...; done'. '[c]' makes the predicate case-insensitive (the unified log writes phrases in mixed
1128 * case across versions).
1129 * **Network listener**: '(echo \"n.add State:/Network/Global/IPv4\"; echo \"n.watch\"; while true; do sleep
1130 * 86400; done) | scutil 2>/dev/null' piped through a 'case' match on 'n.state' / 'SCEventUpdate'.
1131 * **run_recover** checks the marker, debounces via '/tmp/nullexit-watcher.last-recovery' (default 10s
1132 * , configurable via 'NULLEXIT_DEBOUNCE_SECONDS'), and shells out to 'bash recover.sh --post-wake'.
1133 * The 10s debounce prevents a single wake event (which typically emits 2-3 log lines) from triggering
1134 * multiple recoveries.
1135 * 'trap ... TERM INT' 'kill 0' cleanly tears down the process group on launchctl 'bootout' so the '
1136 * sleep 86400' inside the scutil pipe also dies.
1137 * 'wait' blocks indefinitely while both pipelines feed it.
1138 4. 'launchd/com.nullexit.wake-recovery.plist' (LaunchAgent in the user domain runs as the console
1139 * user at every login without needing sudo).
1140 * 'Label: com.nullexit.wake-recovery'
1141 * 'ProgramArguments: [\"/bin/bash\", \"<absolute-path-to-your-nullexit-install>/scripts/watcher.sh\"]'
1142 * 'RunAtLoad: true' starts immediately on 'launchctl load' (no need to log out and back in)
1143 * 'KeepAlive: { SuccessfulExit: false, Crashed: false }' **don't** restart on clean SIGTERM exit (so '
1144 * launchctl bootout' doesn't get stuck in a relaunch loop). Also **don't** restart on hard crash: we
1145 * observed that macOS Sonoma's sleep/wake suspend-resume path accumulates orphan watcher.sh PIDs
1146 * because SIGCONT does not reliably clean up the prior instance's listener grandchildren, and a '
1147 * KeepAlive.Crashed=true'-driven relaunch actively races with the still-live prior process. The script
1148 * enforces single-instance on its own via a 'flock'-style PID-file lock, so a relaunch-on-crash would

```



```

    be skipped by the lock and produce 0 listener coverage until the live instance happened to die.
    Trade-off: a real crash leaves the user without auto-recovery until they run 'launchctl load -w'
    manually acceptable because (a) gateway still works without auto-recovery, (b) a relaunch wouldn't
    fix whatever caused the crash, and (c) the gateway-recovery itself is observably broken (no DNS, no
    exit-node) if it does go down.
1108 * 'ProcessType: Background' hint to App Nap not to suspend us.
1109 * 'StandardOutPath/StandardErrorPath: /tmp/nullexit-watcher.{out,err}.log' separates launchd-level
    diagnostics from the script's own '/tmp/nullexit-watcher.log' (which 'exec >>' redirects).
1110
1111 **Install (one-time).** The plist lives in the repo at '**launchd/com.nullexit.wake-recovery.plist' for
    version control. The installed (live) copy must land in '~/Library/LaunchAgents/' so launchd picks
    it up on the user session.
1112
1113 ```bash
1114 # Copy the plist to the user-launchd location (run from repo root)
1115 cp ./launchd/com.nullexit.wake-recovery.plist \
1116 ~/Library/LaunchAgents/
1117
1118 # Substitute __NULLEXIT_HOME__ with your local install absolute path.
1119 # The plist uses WorkingDirectory + a relative program arg, so this
1120 # single substitution is the only place it references your machine.
1121 sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" \
1122 ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1123
1124 # Validate the XML
1125 plutil -lint ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1126
1127 # Load + persist this user session + every future login
1128 launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1129
1130 # Confirm it actually started
1131 launchctl list | grep nullexit
1132 ```
1133
1134 To **stop** the watcher (without uninstalling):
1135
1136 ```bash
1137 launchctl bootout gui/$UID/com.nullexit.wake-recovery
1138 # or:
1139 launchctl unload -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1140 ```
1141
1142 To **uninstall** completely:
1143
1144 ```bash
1145 launchctl bootout gui/$UID/com.nullexit.wake-recovery 2>/dev/null || true
1146 rm ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1147 ```
1148
1149 **Why this fix was preferred over alternatives.**
1150
1151 * **Polling with 'StartCalendarInterval'** would have given ~30-60s detection latency per wake. The
    unified-log stream is sub-second.
1152 * **Sleepwatcher** would have covered wake events but not network changes. We'd still need a separate '
    scutil' listener, and introducing a third-party dependency for what amounts to ~150 lines of shell
    is unjustified.
1153 * **Forcing 'caffeinate -l'** doesn't exist and the closest equivalent ('caffeinate -s') only works on AC
    power. The whole point of the gateway is to stay usable on battery while traveling lid close must
    NOT silently kill the gateway.
1154 * **A kernel extension** is impossible (no entitlement) and out of scope.
1155
1156 **Reproduction recipe.** Test both events in isolation to confirm the fix actually closes the gap.
1157
1158 * ***(A) Lid-only repro**:: with the gateway up, run in a terminal: 'pmset displaysleepnow && sleep 30 &&
    pmset wake' (without physically closing the lid). Watch the wake wake-fires the watcher post-wake
    routine runs login should still work in <5s after wake. Compare to baseline pre-fix: pre-fix the
    gateway would be dead for 30-90s.
1159 * ***(B) Roam-only repro**:: with the gateway up: 'networksetup -setairportpower off && sleep 30 &&
    networksetup -setairportpower on' (or simply walk out of Wi-Fi range and back). The captive-portal
    WILL repave DHCP DNS for a few seconds; the watcher fires on the second 'State:/Network/Global/IPv4'
    change (after the Wi-Fi re-associates) and the post-wake routine re-hijacks DNS within 2-3s.
1160
1161 **Operative files.**
1162
1163 * 'recover.sh' added arg parsing; wrapped destructive sections behind 'POST_WAKE'; added re-hijack DNS /
    refresh exit-node / force-recreate-if-unhealthy path.
1164 * 'toggle.sh' added 'write_gateway_active_marker' / 'clear_gateway_active_marker' called at the END of
    START and TOP of STOP / cleanup_handler.
1165 * 'scripts/watcher.sh' new long-running daemon.
1166 * 'launchd/com.nullexit.wake-recovery.plist' launchd LaunchAgent descriptor.
1167 * '~/Library/LaunchAgents/com.nullexit.wake-recovery.plist' the installed copy (one-time copy).

```

```

1168
1169 **Honest assessment (read before testing).** This fix has two asymmetric halves:
1170
1171 * **Network / roam / captive-portal path RELIABLE.** 'scutil n.watch' on 'State:/Network/Global/IPv
{4,6}' catches every Wi-Fi roam, captive-portal DHCP-rebind, hotspot handoff, and VPN change with
zero missed events. To validate, drop Wi-Fi for 30 s and re-add it; '/tmp/nullexit-watcher.log' will
record a 'NET:' trigger and 'recover.sh --post-wake' will run within ~10 s of the network coming
back.
1172 * **Lid close / wake path BEST-EFFORT on macOS Sonoma+.** Apple *suspends* LaunchAgents process-state
during forced sleep and *resumes* them on wake; it does NOT re-launch them every wake. As a result:
(a) 'log stream' is a live subscription events emitted during the agent's suppressed window are
DROPPED, not buffered, so the wake event that prompted the resume is invisible to 'log stream' on
resume; (b) the "fire on startup" guard runs once on 'launchctl load -w' and after a 'KeepAlive'
crash-recovery, NOT on every successive wake; (c) 'subsystem == "com.apple.powermanagement"' will
catch FUTURE emissions (display dim/restore, manual sleep/wake) but not the lid-close/open wake
directly.
1173
1174 In practice the lid-wake gap is short, because 'toggle.sh''s existing DNS Watcher already polls every 5 s
and re-hijacks DNS (so DNS recovers within 5 s of any roam or wake), 'tailscaled' self-heals via
DERP fallback within 3060 s, and 'gluetun' reconnects via its healthcheck retry within 30 s. The
user-visible pain on lid open is ~30 s of wonky DNS, not a permanent outage.
1175
1176 **Test the ROAM path FIRST.** If ROAM works in your testing, the lid-wake gap is acceptable. If lid-wake
handling is critical, the two clean follow-ups are:
1177 1. **Sleepwatcher** (one canonical install): 'brew install sleepwatcher'; drop 'recover.sh --post-wake'
into '~/sleep' and '~/wake' so Sleepwatcher invokes it directly on every wake without the launchd
propagation problem.
1178 2. **Move the watcher to a LaunchDaemon** (root) and use 'IOPMSchedulePowerEvent' via a tiny C wrapper
not portable to shell.
1179
1180 A truly reliable shell-only solution to "wake events from inside a LaunchAgent" does not exist on macOS
Sonoma+.
1181
1182 **Empirical lessons from deployment.** Two findings came up while deploying this fix end-to-end;
recording so the next operator doesn't lose an evening to each.
1183
1184 1. **Bash function-call-before-definition gotcha.** While introducing the gateway-active marker helpers,
the defensive 'clear_gateway_active_marker' call was inserted at the very top of 'toggle.sh' (right
after 'set -e'), but the function definition lived near 'PID_FILE="/tmp/nullexit-cafeinate.pid"'
~90 lines down. Bash parses top-to-bottom and resolves names at first invocation, so the call site
exited with code 127 ('clear_gateway_active_marker: command not found'). The gateway stayed
unreachable from the START path even though 'bash -n' passed (the parser doesn't catch order-of-
resolution errors). Rule: define functions BEFORE any block that calls them. Verify with 'grep -n '^
name()\s*{' followed by 'grep -n '^name\s*$' the latter must appear on a line number AFTER the
former.
1185 2. **'networksetup -setairportpower off+on' is not a faithful roam reproduction.** Synthetic Wi-Fi radio-
cycling ('sudo networksetup -setairportpower Wi-Fi off' for 30 s, then back on) is expected to
produce a 'NET:' trigger in 'scutil n.watch' but produced ZERO during testing because '
setairportpower' powers the radio at the driver level while leaving the network SERVICE in the "up
" state from scutil's perspective, so no 'n.state State:/Network/Global/IPv4' event is generated. A
real roam (macOS briefly loses the AP and re-associates into a different one) DOES fire the event
because the link actually changes.
1186 Better reproduction recipes in priority order:
1187 * Walk between two real SSIDs via 'networksetup -switchtolocation' (or actual physical station
movement) guaranteed to fire the scutil event.
1188 * Force a DHCP rebind on the same SSID with 'sudo ipconfig set en0 DHCP' triggers 'State:/Network/
Global/IPv4' to update.
1189 * Trust the existing 30-second DNS Watcher inside 'toggle.sh' for the DNS path: it re-hijacks DNS
automatically. The post-wake watcher adds clear value mostly for tailscale exit-node re-assertion
and warp gluetun force-recreate, both of which self-heal within ~30 s anyway.
1190
1191 #### 15.5.2 July 8, 2026: Network Watcher Event Mismatch, Sudo-in-Background Crash, and Loop-Prevention
1192 **Symptom:**
1193 1. Switching Wi-Fi or waking the Mac from sleep did not trigger a post-wake recovery ('recover.sh --post-
wake') automatically.
1194 2. The network connection would eventually get completely sinkholed/blocked. If the user tried to restart
or wait, it did not self-heal.
1195 3. During the recovery attempt, checking the public IP on the host would return the unencrypted campus/
ISP IP (starting with '132.') instead of the encrypted tunnel IP.
1196
1197 **Root Cause:**
1198 - **Bug 1 (Network Watcher Mismatch):** In 'scripts/watcher.sh', the network change listener watched '
State:/Network/Global/IPv4' changes via 'scutil n.watch' but matched them against '*n.state*|*
SCEventUpdate*'. On macOS, the actual output is 'changedKey [0] = State:/Network/Global/IPv4'.
Because the pattern did not match, all network switches were silently ignored.
1199 - **Bug 2 (Sudo Caching Background Crash):** When the background WARP Watcher eventually detected the
broken tunnel (after 6 failures/30s), it triggered the default nuclear 'recover.sh' (with 'POST_WAKE
=false'). Because there was no active TTY in the background context, the 'sudo -v' caching check
aborted instantly with a terminal required error. Due to 'set -e', 'recover.sh' died immediately
before resetting DNS, disabling the packet-filter killswitch, or stopping containers, leaving the
host network completely sinkholed.

```

```

1200 - **Bug 3 (Post-Wake Infinite Loop):** Once the watcher patterns were fixed, 'recover.sh --post-wake'
    triggered successfully on network changes. However, because the script deleted the default routing
    entries at the start, spent ~15 seconds rebuilding 88 DERP relay bypass routes, and re-added the
    default routes at the end, this execution time exceeded the watcher's 10-second debounce threshold.
    Furthermore, the watcher recorded the *start* timestamp in the debounce file. Consequently, the
    route changes made by the script itself triggered the watcher again immediately after completion,
    trapping the system in an infinite loop of recovery runs. During this loop, the VPN routes were
    constantly deleted, resulting in the host leaking traffic through the real ISP interface (an IP
    starting with '132.') for 15s out of every cycle.
1201
1202 **Resolution:**
1203 - **Fixed Watcher Matching:** Updated 'scripts/watcher.sh' to match '*changedKey*' and '*State:/Network/
    Global/IPv*' to correctly catch network changes.
1204 - **Bypassed Background Sudo:** Added '--auto'/'--non-interactive' flags to 'recover.sh' and added a TTY
    check '[ -t 0 ]' to skip interactive 'sudo -v' authentication when running headlessly in the
    background. Updated the WARP Watcher call in 'toggle.sh' to pass '--auto'.
1205 - **Prevented Infinite Loops:** Modified 'scripts/watcher.sh' to write the **completion** timestamp of '
    recover.sh' to 'DEBOUNCE_FILE' (instead of only the start timestamp). This ensures all buffered
    network config events generated during route reconstruction are evaluated against the end-of-run
    timestamp and successfully debounced, breaking the infinite loop.
1206 - **Centralized & Timestamped Logs:** Redirected 'watcher.sh' stdout/stderr from '/tmp/nullexit-watcher.
    log' to the repository's main 'output.log' and updated 'scripts/common.sh's 'step', 'ok', 'warn', '
    fail', and 'die' helpers to prefix logs with a local '[YYYY-MM-DD HH:MM:SS]' timestamp.
1207
1208 #### 15.5.3 Incident Post-Mortem: Startup Pre-Flight Check Tailscale Race (July 10, 2026)
1209 **Symptom:** Immediately after startup via 'toggle.sh', the pre-flight connectivity check '[1/3] Gateway
    reachable via Tailscale' returned 'FAIL', forcing 'toggle.sh' to skip the full exit-node activation
    and fall back to SOCKS5/local DNS proxy mode.
1210
1211 **Root Cause:** A timing race condition during startup. In 'toggle.sh', the host joins the mesh using '
    tailscale up --reset' (Phase A) right before performing the pre-flight check. Because 'tailscale up'
    resets and reconfigures the host's 'tailscaled' daemon, the local daemon must fetch the latest
    netmap asynchronously from the coordination servers. This network map propagation and route
    establishment takes a few seconds. Running 'tailscale ping' immediately after 'tailscale up' returns
    (with no delay) causes the check to fail because the local daemon has not yet discovered the
    gateway node '100.100.21.8'.
1212
1213 **Fix (July 10, 2026):** Upgraded pre-flight Check 1 in both 'toggle.sh' and 'scripts/toggle-linux.sh' to
    use a 5-attempt retry loop with a 1-second delay between checks. This allows up to 5 seconds for
    local tailnet routing paths to settle, ensuring the gateway is correctly reached and a pong is
    received before deciding whether to enable exit-node routing.
1214
1215 #### 15.5.4 Incident Post-Mortem: Pre-Flight Check False Negatives due to VPN Firewall STUN Blocks (July
    10, 2026)
1216 **Symptom:** Immediately after a cold boot or full restart of the gateway, the pre-flight connectivity
    check '[1/3] Gateway reachable via Tailscale' failed, causing 'toggle.sh' to skip the full exit-node
    configuration and fall back to SOCKS5/local DNS proxy mode. However, manually running 'tailscale
    ping 100.100.21.8' immediately after startup succeeded in 1ms.
1217
1218 **Root Cause:** An asynchronous Tailscale handshake delay caused by firewall restrictions. Inside the
    container network namespace, the 'warp' container (gluetun) implements a strict firewall that blocks
    all outbound UDP traffic to the public internet except to the designated WireGuard endpoint.
    Because of this, the 'tailscale' container's STUN requests to public STUN servers are blocked,
    causing the local daemon to report 'UDP is blocked, trying HTTPS'. Without UDP STUN capability,
    Tailscale is forced to fall back to a slower DERP-assisted handshake (relayed via HTTPS/TCP). This
    negotiation and direct-path punch-through takes roughly 30 to 35 seconds to establish on cold boots
    (after a full 'tailscale up --reset'). Since the pre-flight check only waited 15 seconds (15
    attempts with 1-second sleeps), the check timed out and aborted before the path completed.
1219
1220 **Fix (July 10, 2026):** Increased the pre-flight ping retry limit from 15 to 45 attempts (45 seconds) in
    both 'toggle.sh' and 'scripts/toggle-linux.sh'. This provides sufficient time for the DERP-assisted
    handshake to complete on cold boots, while still passing immediately (in 1-2 seconds) on warm
    starts or once the path is established.
1221
1222 #### 15.5.5 Incident Post-Mortem: Docker Image Build DNS Failures on Immediate Restart (July 10, 2026)
1223 **Symptom:** When executing 'toggle.sh --restart', the script successfully stops the gateway but then
    immediately fails during startup during 'docker compose build' with DNS resolution errors:
1224 'failed to do request: Head "https://registry-1.docker.io/...": dial tcp: lookup registry-1.docker.io on
    192.168.5.1:53: no such host'.
1225
1226 **Root Cause:** A race condition between physical link negotiation and virtual machine startup. The STOP
    path of the restart command invokes 'cleanup_network_state' which power-cycles the host's physical
    Wi-Fi interface ('en0') to flush stale routing states. Immediately after, 'toggle.sh' starts the
    gateway boot sequence. Because 'toggle.sh' did not verify the status of the physical interface, it
    proceeded to boot Colima and build Docker containers while 'en0' was still negotiating a DHCP lease.
    Consequently, the host (and by extension, the VM bridge) had no active internet gateway/DNS path,
    causing Docker's registry check to fail.
1227
1228 **Fix (July 10, 2026):**
1229 1. Created a unified 'wait_for_dhcp_settle' helper function in 'scripts/common.sh' that polls 'ipconfig
    getpacket' and 'ifconfig' until a valid IP and router are assigned. To handle typical macOS Wi-Fi

```

```

1230     reassociation and DHCP negotiation latency (which commonly takes 15-20 seconds after a power-cycle),
        this loop was configured to poll up to 60 times (30 seconds total).
1231 2. Injected a call to 'wait_for_dhcp_settle' at the beginning of the 'START' action in 'toggle.sh' (right
        before checking Colima status) to block container builds until the host's physical network is
        online.
1232 3. Refactored 'recover.sh' to reuse the unified 'wait_for_dhcp_settle' function.
1233 ##### 15.5.6 Watcher.sh Suppressed Exit Code Bug (Fixed)
1234 **Observation:** The 'scripts/watcher.sh' script is a long-running daemon that invokes 'recover.sh --post
        -wake' when it detects system wake or network roam events. Previously, it contained the following
        bug:
1235 ```bash
1236     bash "$RECOVER" --post-wake
1237     local exit_code=$?
1238     ...
1239     return $exit_code || true
1240 ```
1241 The '|| true' mistakenly swallowed the actual '$exit_code' returned by 'recover.sh', causing 'run_recover
        ()' to always return '0' (success), masking any underlying failures in the wake recovery process.
1242
1243 **Fix:** The '|| true' statement was removed from the return line:
1244 ```bash
1245     return $exit_code
1246 ```
1247 Since 'watcher.sh' explicitly omits 'set -e' (to prevent pipe breakages from killing the daemon),
        removing '|| true' safely allows the underlying exit code to propagate without risking daemon
        termination. The watcher daemon was then reloaded ('launchctl kickstart -k') to pick up the script
        changes in memory.
1248
1249 ##### 15.5.7 Network Readiness Race Condition on '--restart' (Fixed)
1250 **Problem:** When running './toggle.sh --restart', the script tears down the gateway and immediately
        begins the startup sequence. On systems with Wi-Fi (or any wireless interface), the OS takes a
        moment to re-associate with the network after the teardown's DNS/routing changes. If the startup
        sequence ran before the OS had a default route, 'get_active_service' would return nothing, routing
        bypass targets would be wrong, and DNS/WARP setup would silently fail resulting in a partially
        configured gateway that self-heals only once the 'watcher.sh' re-runs.
1251
1252 **Root cause:** The original code had no gate before the first network-dependent call ('ACTIVE_SERVICE=$(
        get_active_service)' at the top of the start branch). This was a pure timing race.
1253
1254 **First fix (Wi-Fi specific):** A polling loop on 'ipconfig getifaddr en0' was added to wait for an IPv4
        address on 'en0' before proceeding. This worked but was brittle it only handled Wi-Fi and would
        fail for Ethernet, USB-C adapters, or iPhone USB tethering.
1255
1256 **Generalized fix:** Replaced the 'en0' check with 'route get default | grep 'gateway:'. This detects a
        valid default gateway on *any* interface, covering all connection types. The loop polls every 2
        seconds with a 60-second hard timeout (then proceeds with a warning rather than hanging forever).
        The output prints the detected gateway IP so failures are easy to diagnose in logs.
1257
1258 ```bash
1259 # In toggle.sh, before ACTIVE_SERVICE=$(get_active_service)
1260 while true; do
1261     if route get default 2>/dev/null | grep -q 'gateway:'; then
1262         _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
1263         echo " Network ready (default gateway: $_gw)."
1264         break
1265     fi
1266     ...
1267 done
1268 ```
1269
1270 **Why not ping?** Pinging an IP (e.g. '1.1.1.1') during restart is unreliable because DNS may still be
        hijacked from the previous session, and the exit node routing may not yet be cleared, causing the
        ping to loop back through the tunnel. 'route get default' is a pure local kernel query no packets
        sent.
1271
1272 ### 15.6 WARP Tunnel Liveness & Host Leaks
1273
1274 ##### 15.6.1 Host-Side Traffic Leaking Past the Exit Node (with Remote Clients Routing Correctly)
1275 **Symptom:** 'toggle.sh' reports success. Remote tailnet clients (e.g. an S24 phone) correctly egress
        through Cloudflare WARP and see a Cloudflare IP on 'whatismyip.com'. But on the **HOST Mac running
        the gateway**, 'whatismyip.com' reports the underlying physical ISP's ASN (e.g. 'campus_isp' for a
        university campus Wi-Fi). The gateway container itself is healthy; the leak is specifically on the
        host.
1276
1277 **Root causes:** Three distinct mechanisms can each produce this exact symptom on the host alone, while
        leaving remote clients unaffected. They are mutually rankable so the diagnostic picks the active one
        deterministically.
1278
1279 * ** (A) SOCKS5 fallback lane active. ** 'toggle.sh's three Phase-B pre-flight checks ('tailscale ping'
        AdGuard via 'dig +tcp' WARP internet, toggle.sh ~lines 750-820) sometimes flake during the first

```

```

minute. When ANY of the three fails, the script sets 'SKIP_EXIT_NODE=true' (toggle.sh:849) and
instead activates the SOCKS5 fallback path (toggle.sh:894). The SOCKS5 fallback hijacks DNS to
'127.0.0.1' (so AdGuard filtering still works) but **modern browsers on macOS do NOT honor the
system SOCKS5 proxy** they ignore 'networksetup -setsocksfirewallproxy'. Result: DNS gets ad-
filtered but actual HTTP/S traffic exits direct through the campus ISP. This branch is invisible
from a remote tailnet client because the client's tunnel is unaffected.

1280
1281 * ** (B) IPv6 leak over campus Wi-Fi. ** The Tailscale exit node and WARP tunnel are both IPv4-only (
gluetun + 'post-rules.txt' drop all IPv6 forwarded traffic; see 15.12 IPv6 Exit-Node Leak). But many
university and enterprise APs are dual-stack the host happily accepts AAAA DNS responses and sends
IPv6 traffic straight out 'en0'. macOS happily encrypts that traffic through Tailscale ONLY if
Tailscale has IPv6 advertised; with an IPv4-only exit node, IPv6 traffic skips Tailscale entirely.
This branch is invisible from a phone because iOS/Android Tailscale apps actively sinkhole IPv6 in
their VPN profile, but the macOS host's 'tailscaled' does not.

1282
1283 * ** (C) Route-table freeze and '--accept-routes' reset after 'tailscaled' wake/toggle. ** Running '
tailscale up --reset' clears the exit-node preference and reverts '--accept-routes' back to its
default ('false'). Without explicitly passing '--accept-routes=true', 'tailscaled' will not attempt
to install the default route through the 'utun*' tunnel interface even if the exit node preference
is reconciled. Additionally, macOS occasionally freezes and fails to apply route-table changes (see
15.11 Standalone Tailscale Daemon Freeze). In both cases, 'netstat -rn' shows the default route
still pointing to the physical Wi-Fi gateway (e.g. '172.17.0.1') instead of the Tailscale 'utun*'
interface, causing host traffic to exit direct while remote clients route correctly.

1284
1285 ** Why remote clients don't see this. ** Each remote phone/laptop uses its OWN Tailscale tunnel to reach
the container's '100.x.x.x:443' over the mesh they're end-to-end encrypted regardless of how the
host Mac itself is configured. Only the HOST's own egress sockets (HTTP requests from the host's
browser, curl, etc.) are affected.

1286
1287 ** Resolution. ** A single script: 'scripts/diagnose-host-leak.sh'. It runs the 8 checks, classifies into
one of A / B / C / OK, prints a verdict + a ready-to-run fix command on stdout, AND writes the full
report to 'host-leak-diagnostic-<UTC>.txt' so the user can paste the file back instead of scrolling-
and-copying from the terminal. With '--fix', the matched remediation is applied and the two checks
that could change (warp + ipv6) are re-verified. With '--watch' (or '--watch 30'), it runs the full
baseline diagnostic then enters a continuous loop checking warp/IPv6/default-route every N seconds,
alerting only on state changes logs to 'host-leak-watch.log'.

1288
1289 '' 'bash
1290 # Diagnose only:
1291 bash scripts/diagnose-host-leak.sh
1292
1293 # Diagnose + apply matched fix + re-verify:
1294 bash scripts/diagnose-host-leak.sh --fix
1295 '' '
1296
1297 ** What it prints. ** A 7-section report ('tailscale status', SOCKS5 state, 'cdn-cgi/trace' warp=, IPv6
leak probe, host default route, last toggle.sh output.log hints, resolved gateway Tailscale IP),
followed by a classification block ('Scenario: A | B | C | OK') with the exact command(s) the user
needs to paste into their terminal they don't have to figure out which 'tailscale up' flags match
their environment.

1298
1299 ** What it does NOT do. ** It does not touch 'toggle.sh' or 'toggle.sh's pre-flight logic, does not modify
'.env' (except '--fix' mode appends 'GATEWAY_BYPASS_PING=true' when fixing Scenario A), and does
not restart Docker. It is a read-only diagnostic with an opt-in remediation flag. Safe to run on a
healthy gateway the worst case is a 30-line report that says "OK".

1300
1301 ** Why cdn-cgi/trace over whatismyip.com for the verdict. ** 'whatismyip.com's ISP database routinely
mislabels campus IPv6 ranges as residential ISPs (that's why the local ISP's ASN appears even when
you're routing through Cloudflare). 'cdn-cgi/trace' is hosted by Cloudflare ITSELF and reports 'warp
=on/off' plus the exact edge colo serving the request. It is the only public endpoint that can
truthfully confirm whether the WARP tunnel is in use.

1302
1303 ** Common false alarms. **
1304
1305 * ** 'warp=on' but whatismyip.com still shows the local ISP. ** 'whatismyip.com's ISP-detection database
is unreliable on academic networks. Trust 'cdn-cgi/trace's 'warp=on' over 'whatismyip.com's ISP
string. Use 'https://api.ipify.org' (returns just the IP, no ISP DB lookup) as a third confirmatory
check.
1306 * ** 'tailscale status' shows the host online but exit node preference is missing. ** 'tailscale up --reset
--exit-node=' (empty value) clears any stale preference; 'tailscale up --reset --exit-node=<100.x.x
.x>' re-establishes it. The diagnostic prints the exact command with the resolved TS_IP filled in.
1307 * ** Both IPv6 leak AND route-freeze flags set simultaneously. ** Scenario B outranks Scenario C fixing
IPv6 makes IPv4 the only viable egress, after which the route-freeze usually self-resolves within
~30s (or one more 'toggle.sh' cycle).

1308
1309 ** Deep dive: Host Leak Probe Continuous Sub-Second Egress Monitor ('scripts/host-leak-probe.sh'). **
1310
1311 > ** STATUS: REMOVED (historical). ** 'scripts/host-leak-probe.sh' and the 'HOST_LEAK_PROBE' env flag no
longer exist the file is deleted/untracked and nothing launches it (verified 2026-07-15). Its fail-
closed role was superseded by the ** PF kill-switch ** ('enable_killswitch' in 'common.sh') plus the
in-container ** WARP Watcher ** (15.6.2); the only surviving host-leak tool is the on-demand 'scripts/

```

```

1312 diagnose-host-leak.sh'. The text below is kept for design rationale only treat every present-tense
1313 claim (auto-launch, PID file, 'HOST_LEAK_PROBE=false' toggle) as past tense.
1314
1315 *What it is.* 'scripts/host-leak-probe.sh' is a lightweight background daemon (~1.4 MB RSS, ~01% CPU)
1316 that polls 'https://www.cloudflare.com/cdn-cgi/trace' every 300ms directly from the host via 'curl'
1317 not via 'docker compose exec'. This is a fundamentally different vantage point from the WARP
1318 Watcher (15.6.2): the WARP Watcher asks the WARP *container* whether it sees 'warp=on'; the Host
1319 Leak Probe asks what a real browser request leaving the host's physical NIC looks like. The two
1320 blind spots it covers that the WARP Watcher cannot:
1321
1322 1. **Host-side flash-leaks** a Cloudflare anycast edge re-route, a browser reusing a stale pooled
1323 connection across a tunnel state change, or a split-second routing gap during a Gluetun healthcheck
1324 restart (see 15.12.13 Routing Fix 30-second polling acknowledged 14s window).
1325 2. **Host egress while container is healthy** the container can report 'warp=on' while the host's own
1326 traffic exits direct (the three-scenario leak class documented above in 15.6.1).
1327
1328 *Lifecycle.* Started by 'toggle.sh's START path ('start_host_leak_probe') immediately after '
1329 start_warp_watcher'. Controlled by:
1330
1331 | Signal path | What happens |
1332 |---|---|
1333 | 'toggle.sh' STOP | 'stop_host_leak_probe()' SIGTERM + PID file cleanup |
1334 | 'toggle.sh' 'cleanup_handler' (error/INT/TERM/HUP) | Same |
1335 | 'recover.sh' (default, non '--post-wake') | 6d block kill by PID file |
1336 | 'recover.sh --post-wake' | Left running the gateway is still live |
1337
1338 PID file: '/tmp/nullexit-host-leak-probe.pid'. Toggle with 'HOST_LEAK_PROBE=false' in '.env' to disable
1339 at next gateway start.
1340
1341 *Output format.* All events are written to 'output.log' (repo root) with UTC timestamps. Only state *
1342 changes* are logged the probe is silent during normal healthy operation.
1343
1344 '''
1345 [09:10:26] LEAK warp=off ip=45.23.1.1 prev_warp=on prev_ip=104.28.246.50
1346 [09:10:26] ROTATE warp=on ip=104.28.248.11 prev_ip=104.28.246.50
1347 [09:10:26] HOST-PROBE failed/timeout (count=1)
1348 '''
1349
1350 curl transport errors ('TLS handshake failed', 'Connection refused', etc.) are also appended to 'output.
1351 log' via '2>>output.log' nothing is silenced to '/dev/null'.
1352
1353 *Grep cheatsheet.*
1354
1355 '''bash
1356 grep LEAK output.log # real warp=off events from the host
1357 grep ROTATE output.log # Cloudflare anycast edge rotations (harmless)
1358 grep 'HOST-PROBE' output.log # curl failures / timeouts
1359 '''
1360
1361 *Resource budget.* ~1.4 MB RSS, ~0% CPU at rest, ~3 HTTPS requests/second to Cloudflare. Safe to leave
1362 running for hours. Back off the polling interval ('bash scripts/host-leak-probe.sh 1.0') if you
1363 observe probe-failure pile-ups indicating Cloudflare rate-limiting.
1364
1365 *Detection vs prevention.* This script detects leaks; it does not prevent them. A sub-300ms flash-leak
1366 can still slip between two polls undetected. If 'grep LEAK output.log' shows confirmed leak events,
1367 the next step is a kill-switch (prevention) an 'iptables' rule on the host's 'en0' that drops non-
1368 Tailscale, non-local egress whenever the gateway is active. That is a separate engineering task not
1369 yet implemented.
1370
1371 ##### 15.6.2 In-Flight WARP Tunnel Liveness Monitor & Statistical Auto-Shutdown
1372
1373 **Why.** nullexit's core promise is that your IP always appears as Cloudflare. If the WARP tunnel
1374 silently dies due to a UDP NAT timeout, a Cloudflare edge rotation, a Colima VM clock skew causing
1375 TLS cert rejection, or any of a dozen other failure modes the host silently falls back to egressing
1376 through the physical ISP. The user sees the gateway as "up" (containers running, Tailscale
1377 connected, DNS hijacked) but their real IP is exposed. This is the exact class of failure that Ingo
1378 Blechschmidt documented in his chilling 2024 Linux kernel post-mortem: for over two years, LUKS disk
1379 encryption on suspend was silently broken because a refactored kernel syscall returned success
1380 while doing nothing ([source](https://mathstodon.xyz/@iblech/116769502749142438)). The lesson: **a
1381 security mechanism that fails silently is worse than no mechanism at all** it creates a false sense
1382 of safety that prevents the user from taking corrective action.
1383
1384 **Design principle.** The WARP Watcher ('start_warp_watcher()' in 'toggle.sh') is a background daemon
1385 that polls 'cdn-cgi/trace' every 30 seconds and applies a **statistically calibrated threshold**
1386 before triggering any safety response. This threshold distinguishes transient network blips (which
1387 self-heal within seconds) from genuine outages (which require intervention). A single 'warp=off'
1388 reading is meaningless Docker might be slow, the network might be congested, a container
1389 healthcheck might be restarting. But 6 consecutive off readings (30 continuous seconds of downtime)
1390 has vanishingly low probability of being a false positive: even at an unrealistically high 10% false
1391 -positive rate per poll, P(6 consecutive) = 0.1 0.0001%. In practice, the per-poll false-positive
1392 rate is far lower than 10%, making the real probability astronomically small.

```



```

1356 **What it does.**
1357
1358 1. **Always logs.** Every state transition (onoff, offon) is timestamped to 'output.log'. The user can '
grep WARP output.log' to audit the tunnel's entire lifetime. This is the "never fail silently"
mandate.
1359
1360 2. **Tracks consecutive failures.** A counter increments on each 'off' reading and resets to 0 on any 'on
' reading. This ensures the watcher never fires on intermittent flapping.
1361
1362 3. **Auto-shuts down at threshold.** When the counter reaches 'WARP_FAIL_THRESHOLD' (default 6,
configurable in '.env'), the watcher declares a statistically significant outage and runs 'recover.
sh' the nuclear recovery script that disconnects Tailscale, stops Docker containers, resets DNS to
'1.1.1.1', flushes routes, and power-cycles Wi-Fi. This instantly reverts the host to normal (un-
encrypted) internet, guaranteeing no traffic leaks through a dead tunnel while the user thinks they'
re protected. A trigger file at '/tmp/nullexit-warp-shutdown-triggered' prevents double-firing.
1363
1364 4. **Exits cleanly.** After shutdown, the watcher exits (the gateway is dead, there's nothing left to
monitor). 'stop_warp_watcher()' is also called by 'toggle.sh's STOP path and 'cleanup_handler' to
terminate the daemon cleanly during normal teardown.
1365
1366 **Configuration.** Set 'WARP_FAIL_THRESHOLD=3' in '.env' for a more aggressive 15s timeout, or '
WARP_FAIL_THRESHOLD=12' for a conservative 60s window. The variable is parsed from '.env' using the
same 'grep' pattern as 'GATEWAY_MSS' and 'GATEWAY_BYPASS_PING'.
1367
1368 **Log sample.** After a real outage (or a forced test via 'docker compose pause warp'):
1369
1370 ```
1371 [2026-07-03T21:15:17Z] WARP DOWN cdn-cgi/trace reports warp=off
1372 [2026-07-03T21:15:47Z] WARP SHUTDOWN 6 consecutive failures (threshold=6). Running recover.sh to kill
the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.
1373 [2026-07-03T21:15:52Z] WARP SHUTDOWN recover.sh completed, watcher exiting. Your IP is now your real ISP
IP.
1374 ```
1375
1376 Or, if WARP self-recovers before the threshold:
1377
1378 ```
1379 [2026-07-03T21:15:17Z] WARP DOWN cdn-cgi/trace reports warp=off
1380 [2026-07-03T21:15:32Z] WARP RECOVERED warp=on after 3 consecutive off readings (threshold was 6)
1381 ```
1382
1383 **Relationship to other monitors.** The WARP Watcher is the **innermost safety layer** it runs as a
child of 'toggle.sh' and only while the gateway is active. It complements (does not replace) the '
scripts/watcher.sh' daemon (15.5.1, which handles sleep/wake and Wi-Fi roam) and 'scripts/diagnose-
host-leak.sh' (15.6.1, which is a manual diagnostic tool). The WARP Watcher is the only component
that actively verifies end-to-end tunnel liveness and takes autonomous action when it fails.
1384
1385 #### 15.6.3 Incident: The Overnight Silent IP Leak (July 2026)
1386 **The Problem:** The host leaked traffic over its raw, unencrypted 'eth0' IP continuously for roughly 8
hours, completely bypassing the VPN. The 'toggle.sh' state and the container liveness checks all
reported the gateway as healthy and active.
1387
1388 **The Investigation:**
1389 1. **Sleep/Wake Checks:** We initially suspected macOS sleep cycles broke the routing table. A check of '
pmset -g log' confirmed the laptop literally never slept (0 sleep events recorded since boot),
ruling out all sleep/wake code paths.
1390 2. **Keepalive Expiry:** We discovered 'docker-compose.yml' was missing '
WIREGUARD_PERSISTENT_KEEPLIVE_INTERVAL'. Without a keepalive, standard router NAT tables (which
typically garbage collect UDP after 30 to 120 seconds of silence) quietly expired the WireGuard port
mapping overnight.
1391 3. **Container State Masking:** Because WireGuard is stateless, the 'warp' container didn't immediately
crash. It stayed "Up", which meant Docker's healthchecks and our 'toggle.sh' liveness checks never
noticed the tunnel inside it was totally dead.
1392
1393 **The Solution:**
1394 1. Added 'WIREGUARD_PERSISTENT_KEEPLIVE_INTERVAL=25s' (the WireGuard standard to beat standard 30s NAT
timeouts) to keep the UDP tunnel permanently pinned open through the router.
1395 2. Injected a **Fail-Closed Kill-Switch** into 'routing-fix.sh': A 30-second loop now actively verifies
tunnel health by running 'curl' over 'tun0' to Cloudflare's trace endpoint. If it fails 3
consecutive times, it immediately rewrites the default route to 'blackhole', severing all traffic
instead of letting it leak through 'eth0'.
1396 3. Added a listener in 'watcher.sh' that detects this fail-closed event and instantly pushes a native
macOS UI notification to the desktop, alerting the user to manually intervene.
1397
1398 > [!NOTE]
1399 > **Why dual failure systems?**
1400 > The system now relies on two independent failure handlers that cover each other's blind spots:
1401 > 1. **Automated Self-Healing (WARP Watcher + 'recover.sh')** Triggered when the 'warp' container logs '
warp=off'. This happens when the server actively kicks the client. Because the container *knows* it
is disconnected, it triggers 'recover.sh' to safely reboot Docker and reconnect.

```

```

1402 > 2. **Fail-Closed Kill-Switch ('routing-fix.sh' + 'watcher.sh'): Triggered when a silent network
      failure occurs (e.g., NAT timeouts). The container incorrectly believes it is connected ('warp=on'),
      so auto-recovery never fires. Instead of auto-recovering (which risks falling back to raw Wi-Fi mid
      -process while the user isn't looking), this kill-switch actively blackholes the internet and forces
      a manual intervention via a desktop notification.
1403
1404 ### 15.7 Tailscale P2P, SNAT & Control Plane
1405
1406 #### 15.7.1 macOS SNAT Endpoint Poisoning & The P2P Blackhole (July 10, 2026)
1407 **Symptom:** When a Tailscale client (like an S24) connected to the exact same Wi-Fi network as the macOS
      gateway, it completely lost internet connectivity. However, if the S24 connected to a Windows PC's
      hotspot on the same network, the tunnel worked flawlessly.
1408
1409 **Root Cause:** Claude's critique correctly dismantled the port collision theory: the gateway daemon wasn
      't failing to bind, and macOS wasn't load-balancing ports. The real culprit was **macOS SNAT Source
      Port Mangling poisoning Tailscale's path discovery**.
1410
1411 Here is the exact step-by-step of the "infinite loop" blackhole:
1412 1. **The P2P Handshake:** The S24 and the Mac are on the same Wi-Fi ('192.168.1.x'). The S24 discovers
      the Mac's local IP and sends a UDP hole-punch packet to '192.168.1.5:41641'. Colima successfully
      forwards this to the gateway container.
1413 2. **The Bypass Rule:** The gateway replies. Because 'routing-fix.sh' forces all Tailscale UDP traffic
      out 'eth0' to bypass WARP, the reply goes to the macOS host.
1414 3. **macOS SNAT Mangling:** macOS routes the reply out 'en0' to the S24. Because the packet crosses from
      the Colima VM bridge to the Wi-Fi interface, macOS applies Source NAT (SNAT). Critically, macOS **
      randomizes the source port** (e.g., changes it from '41641' to '58291').
1415 4. **Endpoint Poisoning:** The S24 receives the authenticated Tailscale packet from '192.168.1.5:58291'.
      Tailscale's 'magicsock' is designed to handle NAT dynamically, so it assumes the Mac has roamed to
      port '58291'! The S24 updates its endpoint and starts sending all its encrypted internet traffic to
      '192.168.1.5:58291'.
1416 5. **The Blackhole:** macOS has no port forward for '58291'. It drops 100% of the S24's outbound data
      packets. The connection is completely dead.
1417
1418 **Why did the Hotspot work?** When the S24 connected to the PC hotspot, it was behind the PC's NAT. The
      S24 sent the packet, PC NATted it, and the Mac replied (mangled to '58291'). When the reply hit the
      PC, the PC's strict NAT dropped it because the source port ('58291') didn't match the destination
      port it originally sent to ('41641'). Because the mangled packet was dropped, the S24 never poisoned
      its endpoint! It safely gave up on P2P, fell back to the 100% reliable TCP DERP relay, and the
      internet worked perfectly.
1419
1420 **Fix & Resolution (July 10, 2026):**
1421
1422 **Phase 1 Port disambiguation:** Changed the gateway container's Tailscale listen port from the default
      '41641' to '41642' across 'docker-compose.yml', 'post-rules.txt', and 'routing-fix.sh'. This
      prevents any bind conflict with the macOS host's native Tailscale daemon.
1423
1424 **Phase 2 RFC1918 DROP rules:** Updated 'routing-fix.sh's 'add_tailscale_p2p_bypass()' to explicitly
      drop all outgoing Tailscale UDP packets destined for private subnets ('192.168.0.0/16',
      '10.0.0.0/8', '172.16.0.0/12') when 'TAILSCALE_ALLOW_LAN_P2P=false'. This surgically kills the local
      P2P hole-punch attempt *before* it can trigger the macOS gvproxy SNAT mangling. The S24 receives a
      clean failure and immediately locks onto the public DERP relay with 0% packet loss.
1425
1426 **Phase 3 Automatic network detection:** Because 'TAILSCALE_ALLOW_LAN_P2P=false' breaks P2P on trusted
      hotspots that genuinely don't have AP isolation, a '.env' toggle and a full auto-detection system
      were implemented:
1427
1428 - **'.env' toggle:** 'TAILSCALE_ALLOW_LAN_P2P=true/false' static fallback, used only if auto-detection
      cannot run.
1429 - **'watcher.sh' auto-detection ('detect_lan_p2p_mode'): On every network change, 'watcher.sh' runs two
      checks:
1430 1. **WPA2-Enterprise detection:** 'airport -I' is used to read the current Wi-Fi 'link auth' field. If
      the auth type does not contain '-psk' (e.g., it is plain 'wpa2' = 802.1x), the network is classified
      as enterprise and P2P is forced off. *(Note: the 'airport' binary was removed in macOS Sonoma 14.4
      the detection now uses 'system_profiler SPAirPortDataType'; see 15.11 for the full migration.)*
1431 2. **AP Isolation probe:** For WPA2-Personal networks, a short 'ping' is sent to the default gateway.
      If the gateway responds (it always should on a real home/hotspot network), P2P is allowed. If not,
      AP isolation is inferred and P2P is forced off.
1432 - The result ('true' or 'false') is written to '.lan_p2p_detected' in the repo root.
1433 - **'routing-fix.sh' dynamic reading:** 'add_tailscale_p2p_bypass()' re-reads '.lan_p2p_detected' on
      every 30-second loop iteration and atomically adds or removes the RFC1918 DROP rules. **No container
      restart is needed when switching networks.**
1434
1435 **Known Unknowns (Pending Verification):**
1436 - The gvproxy SNAT source port mangling theory is mechanistically sound and backed by observed behavior (
      gvproxy's UDP proxy changelog has documented edge cases in this exact code path), but has not yet
      been confirmed via 'tcpdump -i en0 udp portrange 40000-60000' while triggering a P2P handshake from
      the phone. A packet capture cross-referenced with 'tailscale ping' endpoint reports on the phone
      side would confirm the theory definitively.
1437 - Claude's recommendation: run 'tcpdump -i en0 udp port 41642 or portrange 40000-60000' on the Mac while
      triggering the handshake from the phone to see whether the reply leaves with a different source port
      than '41642'.

```

1438 - ****Client-Side VPN Cycling (Roaming Recovery):**** An observation (July 10, 2026) suggests that when a hung DNS probe state is triggered by roaming between Access Points (APs) on a WPA2-Enterprise network, simply toggling the VPN connection (Tailscale) off and on locally on the client phone immediately resolves the issue. You do not need to cycle the phone's physical Wi-Fi. This reinforces the theory that roaming on Enterprise networks leaves a stale/poisoned P2P endpoint in the phone's Tailscale magicsock state, which gets instantly flushed upon a local VPN restart. (Status: Possible but not formally verified).

1439

1440 **#### 15.7.2 July 6, 2026: Packet Filter (pfctl) Defaulting Local LAN Rules to TCP-Only**

1441 ****Symptom:**** Devices on the exact same local Wi-Fi network as the gateway were experiencing ~500ms latency to the gateway instead of the expected ~80ms.

1442

1443 ****Root Cause:**** When writing the 'pf.conf' rules to whitelist local LAN traffic ('pass out quick on \$ext_if to { 192.168.0.0/16, ... }'), the rule omitted an explicit protocol. The macOS 'pfctl' compiler implicitly applies state tracking to all 'pass' rules. However, because state tracking ('keep state') natively evaluates TCP sessions, 'pfctl' automatically appended the 'flags S/SA' modifier to the compiled rule in the kernel. This silently restricted the whitelist to ****TCP only****, causing all local UDP traffic to be dropped by the default-deny kill-switch. Since Tailscale relies on UDP port 41641 for direct P2P connections, the local devices were unable to hole-punch and were forced to fall back to routing traffic through a remote Tailscale DERP relay, massively inflating latency.

1444

1445 ****Resolution:**** Split the single implicit rule into explicit 'proto tcp', 'proto udp', and 'proto icmp' rules in 'scripts/pf.conf' to guarantee the macOS compiler allows all three core transport protocols on the local subnet without mutating the rule into a TCP-only lock.

1446

1447 **#### 15.7.3 Incident Post-Mortem: PF Kill-Switch Bricks Host Internet in SOCKS5 Fallback Mode (July 10, 2026)**

1448 ****Symptom:**** When the pre-flight checks fail (e.g. because host Tailscale was unauthenticated/logged out), the script falls back to SOCKS5 proxy mode but the host immediately loses all internet connectivity and tailscaled reports 'You are logged out... connect: no route to host'. Even running 'sudo tailscale up' to log in fails because the connection to the control plane times out.

1449

1450 ****Root Cause:**** An architectural mismatch. The macOS Packet Filter (PF) Kill-Switch works at the IP level : it blocks all outbound traffic on the physical interface ('en0') except to explicitly whitelisted VPN endpoints, expecting all other traffic to go through the virtual VPN interface ('utun*'). In SOCKS5 fallback mode, the exit-node ('utun*') is NOT enabled; instead, only specific applications (like browsers) use the localhost SOCKS5 proxy, while all other host traffic (including the 'tailscaled' daemon's UDP packets) continues to go through 'en0'. Because the script still enabled the PF Kill-Switch during SOCKS5 fallback, it blocked all non-SOCKS5 traffic on 'en0', completely bricking the host's internet and locking the Tailscale daemon out of the control plane.

1451

1452 ****Fix (July 10, 2026):**** Modified 'toggle.sh' to remove the 'enable_killswitch' call from the SOCKS5 fallback path. The firewall kill-switch is now strictly reserved for exit-node VPN mode, ensuring SOCKS5 fallback leaves the host's direct internet route open for system daemons to authenticate.

1453

1454 **#### 15.7.4 Incident Post-Mortem: Host Tailscale Logouts due to Aggressive reset Flags (July 10, 2026)**

1455 ****Symptom:**** Every time the user runs 'toggle.sh' or 'recover.sh' which disconnects/reconnects the host client, the host client randomly gets logged out, forcing them to re-run 'sudo tailscale up' to authenticate.

1456

1457 ****Root Cause:**** An overly aggressive configuration reset. The host-side 'tailscale up' calls (used in 'disconnect_tailscale_host' and the startup/enable exit node paths) all passed the '--reset' flag. On macOS, '--reset' resets the entire configuration state and authentication token cache of the daemon. Since the script does not pass a 'TS_AUTHKEY' to the host's commands (which are meant for the user's private interactive client), the '--reset' flag effectively logged the user out of their tailnet, requiring interactive re-login.

1458

1459 ****Fix (July 10, 2026):**** Removed the '--reset' flag from all host-side 'tailscale up' invocations in 'toggle.sh' and 'scripts/common.sh'. This ensures that exit-node state transitions (enabling/disabling) are handled safely while preserving the host client's authenticated session.

1460

1461 **#### 15.7.5 Incident Post-Mortem: macOS Tailscale Flag Requirement Abort (July 10, 2026)**

1462 ****Symptom:**** When the 'toggle.sh' stop trap or the roaming watcher triggered, the gateway failed to shut down or disconnect properly. The 'output.log' showed the error: 'Error: changing settings via 'tailscale up' requires mentioning all non-default flags. To proceed, either re-run your command with --reset...'.

1463

1464 ****Root Cause:**** In a previous attempt to stop host-side logouts (15.7.4), we removed the '--reset' flag from all 'tailscale up' invocations. However, if the Tailscale daemon on macOS is already running with non-default flags (like '--ssh' or '--accept-routes'), invoking 'tailscale up' with only a subset of flags (like '--exit-node=') triggers a strict safety check in the Tailscale CLI that aborts the command completely. Because the command aborted, exit-node routing remained stuck.

1465

1466 ****Fix (July 10, 2026):**** Migrated all specific preference changes in 'toggle.sh' and 'scripts/common.sh' from 'tailscale up' to a two-step sequence:

1467 1. 'tailscale up' (with no flags) to safely ensure the interface is brought up using the current authenticated state.

1468 2. 'tailscale set --exit-node=...' to modify the specific target preference cleanly without triggering the missing flags error or requiring '--reset'.

1469

```

1470 ### 15.8 Docker / Colima Lifecycle
1471
1472 ##### 15.8.1 Colima VM Memory / Swap Thrashing Latency
1473 **Symptom:** After running nullexit flawlessly for over 24 hours straight to test stability, the 0.5GB VM
RAM configuration began experiencing severe network latency on the Tailscale mesh later in the day.
1474
1475 **Root Cause:** Services (like AdGuard, Tailscale, Docker logs) slowly accumulate cache and state over
extended periods. Under the tight 512MB physical RAM limit, this slow creep forced the Linux kernel
to aggressively swap memory to the 512MB SSD swap file. The constant disk I/O thrashing blocked
process execution, causing DNS resolutions and packet forwarding to delay significantly (making the
internet feel "very slow").
1476
1477 **Proof (Empirical Observation):** While in this OOM-thrashing state, direct DNS queries through the
Tailscale mesh ('dig @100.100.21.8 google.com') would intermittently hang or take several seconds to
resolve. After restarting with the 600MB configuration, the exact same query immediately dropped to
a lightning-fast **41ms** response time. (Note: Querying the mapped host port via '127.0.0.1:5354'
will always time out due to Gluetun's strict leak-prevention firewall blocking non-VPN ingress
traffic).
1478
1479 **Fix:** Increased VM base physical RAM to 600MB ('--memory 0.6') and reduced the swap file to 400MB.
This provides enough native RAM headroom for long-running state caches without forcing heavy I/O
swapping. Subdomain deduplication still reduces blocklists by ~60%. Memory profiles ('light'/'medium'
/'heavy') let users trade off coverage for memory.
1480
1481 ##### 15.8.2 Missing iptables in routing-fix Container
1482 **Root Cause:** 'alpine:3.20' does not have iptables installed. Commands failed silently with '2>/dev/
null || true'.
1483
1484 **Fix:** Updated 'docker-compose.yml' to 'apk add --no-cache iptables iproute2' before 'scripts/routing-
fix.sh'.
1485
1486 ##### 15.8.3 DOCKER_GW Variable Parsing Bug
1487 **Root Cause:** 'ip route show default' printed two lines; 'awk '{print $3}'' picked up both, causing
route commands to fail.
1488
1489 **Fix:** Added 'head -1' to the parsing chain.
1490
1491 ##### 15.8.4 Hard Reboots Leave Stale Docker Socket in Colima VM
1492 **Problem:** If the macOS host machine is subjected to a hard reboot, sudden power loss, or a kernel
panic, the Colima VM running in the background is abruptly killed. Upon next boot, running 'toggle.
sh' resulted in the contradictory output:
1493 'Colima is already running.' followed by 'Cannot connect to the Docker daemon at unix:///colima/default
/docker.sock.'
1494
1495 **Root Cause:** The hard crash prevents the inner Docker daemon from executing its graceful shutdown
routines. It leaves behind a corrupted 'docker.sock' file inside the VM. On boot, the outer VM spins
up (hence "already running"), but the inner Docker daemon sees the stale socket, assumes another
instance is running, and crashes.
1496
1497 **Fix:** Added an Auto-Healing routine directly into 'toggle.sh' (Step 4b). The script now explicitly
tests the Docker daemon with 'docker ps'. If the daemon is unresponsive despite Colima reporting as
active, it assumes a stale socket and automatically runs 'colima restart' in the background to
cleanly rebuild the VM and socket before proceeding.
1498
1499 ##### 15.8.5 Cloudflare WARP 24-Second Startup Delay (UDP Session Rate-Limiting)
1500 **Context:** When running 'toggle.sh' to turn the gateway ON, 'docker compose up -d' often hangs for
exactly 24-25 seconds waiting for the 'warp' container to become 'healthy'. Users might mistakenly
blame the Go 'rule-compiler' processing 400k+ rules for this delay, as Docker prints 'rule-compiler
Exited 24.3s' at the end of the wait.
1501 **Root Cause:** The 'rule-compiler' actually finishes its work in <2 seconds. The 24-second blockage is
entirely due to Cloudflare WARP's edge server rate-limiting. WireGuard is a stateless protocol with
no "disconnect" packet. When the toggle is turned OFF, the old UDP session state remains active on
Cloudflare's backend for 20-30 seconds. When the toggle is instantly turned back ON, Docker starts a
new WireGuard connection from a new ephemeral UDP source port but uses the **same static
cryptographic key** (from '.env'). Cloudflare detects this as a potential replay attack or key-
sharing abuse, and intentionally drops all packets (rate-limits) for the new connection until the
old session's ghost state fully expires (~20-25 seconds).
1502 **Result:** Gluetun's healthcheck fails repeatedly every 2 seconds until Cloudflare lifts the penalty box
, at which point traffic flows and Docker unblocks the rest of the startup sequence. This is a known
, expected behavior of reusing static keys rapidly on Cloudflare's network and cannot be bypassed.
1503
1504 ##### 15.8.6 July 4, 2026: Multi-stage Docker Builds for Go Ports & Teardown Hang Fix
1505 **Symptom:** The Python implementations of 'logger.py' and 'sync-rules.py' were slow and required a heavy
python runtime inside Alpine containers. Also, 'routing-fix' container teardown was hanging for 30s
during 'toggle.sh' shutdown.
1506 **Root Cause:**
1507 - Python is slower than Go and installing it dynamically via 'apk add' added overhead and complexity to
the containers.
1508 - Docker containers ignore 'SIGTERM' if the init process doesn't handle it, causing a full 30s timeout on
shutdown.
1509 **Resolution:**

```

1510 - Ported 'logger' and 'sync-rules' to Go using Goroutines for massive concurrent speedups.
1511 - Implemented **Multi-stage Docker builds** (using 'golang:1.22-alpine' as builder) to compile the static binaries inside Docker. The final containers are raw Alpine with zero dependencies.
1512 - Added '--build' to 'docker compose up -d' in 'toggle.sh' to ensure updates are consistently applied on boot.
1513 - Added 'stop_grace_period: 1s' to 'routing-fix' in 'docker-compose.yml' which eliminates the 30-second teardown hang instantly.
1514 - Implemented a cryptographic integrity checker ('scripts/crypto.sh') using HMAC-SHA256 and 'NULLEXIT_SEED' in '.env' to prevent tampering of bash scripts. (See also 15.9.4.)

1515
1516 **#### 15.8.7 Colima 'start' Post-Provision Hang (~2 min) & the Docker-Readiness Poll (July 14, 2026)**
1517 ****Symptom:**** 'toggle.sh' took ~197s of wall-clock time to bring the gateway up, and almost the entire delay lived in a single step: 'colima start'. Timestamps in 'output.log' showed the VM reaching 'READY' and creating the 'colima' Docker context in ~8-12 seconds, after which 'colima start' simply *did not return* for ~110 more seconds until the 'run_with_timeout 120' watchdog 'SIGKILL'ed it. The rest of the boot (containers, routing) then proceeded normally.

1518
1519 ****Root Cause:**** A colima-internal hang that occurs *after* 'Successfully created context colima' is printed. It is **mode-independent**. It was initially misattributed to '--network-address' / '--network-mode bridged' needing the 'socket_vmnet' helper, but a controlled test with plain user-mode networking ('colima start --memory 0.6 --vm-type vz', no '--network-address') reproduced the identical ~120s hang (04:40:34 04:42:34, watchdog-killed) while Docker was fully usable the entire time. Crucially, killing the hung 'colima start' process does **not** stop the VM and does **not** corrupt state 'colima status' correctly reports 'running' afterwards and 'docker' keeps working via the 'colima' context. The ~110s was therefore pure dead time waiting on a CLI that had already done its real work.

1520
1521 ****Fix:**** 'scripts/common.sh' now provides 'colima_start_until_ready()'. It launches 'colima start "\$@"' in the background and polls 'docker --context colima info' every 3s. The moment Docker answers (~12s in practice) it stops waiting, reaps the (usually still-hung) starter with 'kill', runs 'docker context use colima' to guarantee the default context is set, and returns 0. If Docker never answers within a 90s cap it returns non-zero and 'toggle.sh' continues to the existing Step 4b 'docker ps' auto-heal (15.8.4). Both Colima start call sites in 'toggle.sh' the LAN-P2P bridged/shared path and the user-mode fast path now go through this helper instead of 'run_with_timeout 120 colima start'.

1522
1523 ****Result:**** The Colima phase dropped from ~120s to ~12s and total 'toggle.sh --restart' wall time from ~197s to ~106s, with the double-tunnel and direct hostcontainer P2P confirmed intact after the change. The key insight: treat "Docker is reachable" not "the 'colima start' CLI returned" as the definition of *started*.

1524
1525 **#### 15.8.8 Toggle-On Startup Optimizations: Rule-Compiler Off the Critical Path (July 14, 2026)**
1526 ****Context:**** A genuine cold toggle-on (VM bounced for RAM, WARP down) measured **995s**. Profiled per-phase with the now-fixed 'DEBUG_TRACE' (15.11.8-A): Colima cold start **19s** (near floor), 'docker compose up' **45s**, tailscale connect+poll **21s**, final checks **7s**.

1527
1528 ****Two wrong hypotheses first (documented so they aren't re-tried):**** (1) that BuildKit '--build' was the ~30s cost it wasn't; every layer was already 'CACHED', so gating '--build' saved only ~2s. (2) that 'docker compose exec' was slow in the poll loop it isn't (0.18s steady-state vs 0.06s for 'docker exec'); the ~5s/iteration was 'tailscale status' **blocking** during cold connect and hitting the 'run_with_timeout 5' cap.

1529
1530 ****Real bottleneck (evidence-backed):**** the 45s 'compose up' was **not** WARP. On a genuine cold start WARP goes healthy fast (tor starts right with it no 15.8.5 same-key penalty; that only bites rapid offon). The long pole is the **rule-compiler** re-deduping ~340k rules inside the 600MB VM (~28s), every toggle **'adguardhome'** started 28s after warp because it waited on 'rule-compiler: service_completed_successfully', even though all 16 remote lists loaded from cache unchanged.

1531
1532 ****Fix** hash-gated, off the critical path (kept in-container):
1533 - 'rule-compiler' is now a **compiled-profiled** service, excluded from 'docker compose up'. 'adguardhome' and 'routing-fix' no longer 'depends_on' it they boot instantly on the 'compiled_rules.txt' / 'ip_blocklist.ipset' already in './adguard/work'.
1534 - 'toggle.sh' runs it on-demand via 'docker compose run --build --rm rule-compiler', **gated** on 'compute_rules_hash()' (** = sha256 of 'black_list.txt' + 'white_list.txt' (the lists the user controls)). Recompile only when that hash changes, when an artifact is missing, or when the compiled file is older than 'RULES_MAX_AGE_DAYS' (default 7, so the remote blocklists still refresh occasionally). A compile failure is **non-fatal** we keep the previous compiled rules rather than bricking the tunnel (ad-block the kill-switch).
1535 - Also trimmed the tailscale-connect poll timeouts (status '5s3s', 'ip -4' '10s5s') to reduce overshoot.

1536
1537 ****Why NOT** run the compiler natively on the Mac host (the tempting idea): **rejected** on purpose. There is no Go on the host, so "native" means either 'brew install go' (a host toolchain dependency that drifts the project from *portable+secure-on-any-device* toward *faster-on-mine*) or committing a prebuilt arch-specific **binary blob** you'd then have to trust/sign. Worse, it would make the **host** write into the './adguard/work' bind-mount violating the deliberate **"single allowed writer is the rule-compiler container"** invariant (hostcontainer UID/perms mismatches corrupt AdGuard's BoltDB store). Staying in-container preserves both properties; the gate delivers the speed.

1538
1539 ****Result:**** unchanged toggle 'compose up' **45s 9s**, total '--restart' **95106s 63s**, with AdGuard still serving the full 342,519 rules (it reuses the existing compiled file) and the double-tunnel + direct hostcontainer P2P intact. Editing 'black_list.txt'/'white_list.txt' flips the hash one-off recompile new rules take effect. Both lists were added to the 'crypto.sh' signed set (they're


```

security inputs), so editing them requires a re-sign which is the same moment you'd want the
recompile. Local state files '.build_hash' / '.rules_hash' are gitignored.
1540
1541 ### 15.9 Permissions, Crypto & Security Hardening
1542
1543 #### 15.9.1 Sudo Credential Caching Removed
1544 **Problem:** The script used 'sudo -v' at startup to cache sudo credentials, plus a background '
SUDO_KEEPER_PID' loop refreshing them every 60 seconds. This required interactive password entry
even with NOPASSWD sudoers configured, because 'sudo -v' itself prompts.
1545
1546 **Fix:** Removed the entire 'sudo -v' + 'SUDO_KEEPER_PID' section. All privileged commands now use 'sudo
-n' (non-interactive), which works silently because the user has NOPASSWD rules in '/etc/sudoers.d/
nullexit' for the specific commands needed ('networksetup', 'python3', 'dscacheutil', 'killall', '
pskill', 'route', 'ifconfig').
1547
1548 #### 15.9.2 Background Sudo Failures
1549 **Context:** The '--post-wake' script is intentionally designed to skip the 'sudo -v' interactive prompt
because it is launched by 'watcher.sh' running via 'launchd' in the background (which has no TTY and
cannot accept password input). Thus, it relies entirely on 'sudo -n' (non-interactive sudo) for all
its internal routing changes.
1550 **Problem:** If the user has not configured the '/etc/sudoers.d/nullexit' file properly, any automated
background wake event will silently fail: 'sudo -n' commands drop out, the WARP bypass routes fail
to apply, and the 'warp' container loses internet connectivity.
1551
1552 #### 15.9.3 Tailscale Hijacking the Default Route (PHYSICAL_GW Bug)
1553 **Problem:** When '--post-wake' runs, it clears the routing table using 'sudo route -n flush' to ensure a
clean slate after the Wi-Fi card roams or wakes up. Because it skips bouncing the Wi-Fi interface (
to make the reconnect faster), macOS's DHCP client does not automatically rebuild the physical
default route. To counteract this, the script must manually restore the physical default route.
Originally, the script tried to capture the physical gateway via 'route get default'. However,
because the gateway is already running and Tailscale is active, the default route is often already
pointing to 'utunX'. When this happens, 'route get default' does not output a 'gateway:' field (it
only outputs the 'interface:'), causing the captured 'PHYSICAL_GW' variable to be empty.
1554 When 'PHYSICAL_GW' is empty, the physical default route is never restored, causing the 'en0' interface to
be isolated from the physical network. The WARP bypass routes then fall back to targeting '-
interface en0' (instead of routing via the gateway), which breaks WireGuard's remote connections
completely.
1555 **Fix:** The script now bypasses the OS routing table entirely and directly queries the hardware DHCP
lease ('ipconfig getpacket en0') to reliably extract the physical router IP, ensuring the physical
route is correctly restored even when Tailscale is active.
1556
1557 #### 15.9.4 July 4, 2026: Consolidation of Core Bash Scripts (Refactoring)
1558 **Symptom:** Over 200 lines of bash logic were directly duplicated across 'toggle.sh' and 'recover.sh' (e
.g. 'restart_tailscaled_daemon', 'stop_sleep_prevention', WARP bypass routes). This caused drift
when one script was updated without the other.
1559 **Root Cause:** Historical separation of responsibilities allowed 'recover.sh' to grow complex alongside
'toggle.sh'.
1560 **Resolution:**
1561 - Massively refactored both scripts by extracting all duplicated networking logic, PID lifecycle
management, and environmental configuration down to 'scripts/common.sh'.
1562 - Specifically, the 'stop_sleep_prevention', 'stop_dns_watcher', 'stop_host_leak_probe', and '
stop_warp_watcher' functions were collapsed into a single 'stop_pidfile_daemon()' abstraction.
1563 - Proxy disable loops were abstracted to 'disable_all_proxies()'.
1564 - The '162.159.192.1/193.1' WARP edge IPs are now dynamically resolved from '.env' instead of hardcoded.
1565
1566 #### 15.9.5 Threat Model: Local Proxy Discovery
1567 The Tor SOCKS proxy and Honey-Port Tripwire only bind to '127.0.0.1' (loopback). The real security
boundary here is "can an unprivileged local process reach the loopback interface," not "does the
malware know the port number."
1568
1569 The port-randomization and the Honey-Port trap only catch blind or generic port-scanners that do not
already know where to look. They **do not protect** against a process that directly reads the '/tmp/
nullexit-ports.env' file, greps the 'toggle.sh' memory space, or runs 'lsof -i' to inspect open
socket bindings.
1570
1571 Because modifying file ownership on '/tmp/nullexit-ports.env' via 'chown'/'chmod' to restrict read-access
introduces an unavoidable Time-Of-Check to Time-Of-Use (TOCTOU) race condition (the file is created
user-owned before privilege escalation, meaning file descriptors can be snatched), it is
deliberately left as a standard user-owned file.
1572
1573 The actual, proper mitigation for preventing a malicious local process from reaching the proxy is not
hiding the port it is running a host-level outbound firewall with strict localhost/loopback filtering
enabled (e.g., LuLu 4.3.0+ or Little Snitch). These firewalls gate access by cryptographic process
identity (binary signature) at the kernel/network-extension level, rather than by port secrecy.
1574
1575 **Verifying Localhost Filtering (The Rogue Binary Test)** To empirically validate that the host firewall
properly intercepts local proxy hijacking, you can compile and execute a completely unknown, un-
whitelisted "rogue" binary to attempt a connection to the proxy port:
1576
1577 ```c
1578 // lulu-test.c

```



```

1579 #include <stdio.h>
1580 #include <stdlib.h>
1581 #include <arpa/inet.h>
1582
1583 int main(int argc, char *argv[]) {
1584     int port = atoi(argv[1]);
1585     int sock = socket(AF_INET, SOCK_STREAM, 0);
1586     struct sockaddr_in server = { .sin_family = AF_INET, .sin_port = htons(port) };
1587     server.sin_addr.s_addr = inet_addr("127.0.0.1");
1588
1589     printf("Rogue binary attempting connection to 127.0.0.1:%d...\n", port);
1590     if (connect(sock, (struct sockaddr *)&server, sizeof(server)) < 0) {
1591         printf("Connection BLOCKED by firewall.\n");
1592         return 1;
1593     }
1594     printf("Connection SUCCESSFUL (Firewall bypassed!).\n");
1595     return 0;
1596 }
1597 '''
1598 Compile with 'gcc -o lulu-test lulu-test.c'. When executed against the '$TOR SOCKS_PORT', a properly
    configured host firewall (with "Allow Loopback" disabled in LuLu Preferences) will immediately pause
    the socket execution and generate a desktop alert for the unrecognized binary signature. This
    physically stops targeted local malware from exfiltrating data through the Tor tunnel, proving the
    threat model mitigation.
1599
1600 ##### 15.9.6 Unlocking Mode-Locked Output Files Without 'chmod' ( 'unlock-files.sh' )
1601 **The one-command fix:** 'bash scripts/unlock-files.sh' replaces every known locked inode in the project
    with a fresh writable one. No 'chmod'. Covers '.env' (mode '000'), 'adguard/work/userfilters/
    compiled_rules.txt', 'ip_blocklist.ipset', 'cache/*.txt', 'cache/ip/*.txt', and 'adguard/work/data/
    filters/*.txt' (mode '0444'). Run it once after setup or after pulling an old repo; 'toggle.sh' and
    'sync-rules.py' will then work without permission errors.
1602
1603 **Context:** The nullexit project has a strict project-wide policy of 'no chmod from scripts' (see README
    6). This rule exists because every prior 'chmod 0444' (post-write tamper-proof lock) and 'chmod
    000' (post-toggle '.env' lock) call from 'scripts/sync-rules.py' / 'toggle.sh' could leave a file
    permanently unreadable on disk if the script exited early, was killed mid-flight, or simply set a
    restrictive mode without ever being re-flipped. We've stripped every 'chmod' from the codebase. But
    files **written by older versions of those scripts** are still on disk in those restrictive modes
    ('0444' for 'compiled_rules.txt', 'ip_blocklist.ipset', 'data/filters/<id>.txt'; '000' for '.env'
    after end-of-toggle lock). Re-running 'toggle.sh' or 'scripts/sync-rules.py' against those leftovers
    hits 'PermissionError: [Errno 13] Permission denied' immediately.
1604
1605 **Why the stale permissions also block 'mv'/'rm' of the parent folder:** The restrictive modes ('000' on
    '.env', '0444' on output files) don't just block writes they prevent the kernel from unlinking the
    file's dentry during a 'mv' or 'rm' of the **containing directory** ('adguard/work/'). macOS
    enforces this at the VFS layer: you own the directory, but you can't delete a directory that
    contains entries you can't write to. This made the entire 'adguard/work/' tree unmovable/undeletable
    until the locked inodes inside it were replaced. 'unlock-files.sh' resolves this permanently by
    swapping each locked inode for a fresh '0644' one via atomic rename ('cp' + 'mv'), which operates on
    the directory entry rather than the file contents no 'chmod' called, no permission bypass needed.
1606
1607 **Problem:** You (the owner) cannot 'cat', 'cp', 'rm', '>' redirect, or any normal POSIX write against a
    'mode=000' file even though you own it the basic mode bits block ALL access for owner, group, and
    other. The script does not call 'chmod' and you want a permanent fix that never relies on a chmod
    from any future run either.
1608
1609 **Solution:** Replace the locked inode with a fresh readable one. 'mv' is atomic; mode bits live in the
    inode, not the directory entry. Two flavors cover every case we hit on a real machine:
1610
1611 **Flavor A locked file is mode '000' (owner can't even READ it):**
1612 NOPASSWD sudoers ('/etc/sudoers.d/nullexit') already includes '/usr/bin/python3'. So we read via 'sudo
    python3' (which has the DAC override root has) and pipe the bytes around as base64 in user space,
    where the redirect happens with the user's normal umask = '0644':
1613 '''bash
1614 sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb')).
    read()).decode()" > /tmp/snap.b64
1615 base64 -d < /tmp/snap.b64 > <FILE>.new
1616 mv -f <FILE>.new <FILE>
1617 ls -la <FILE> # mode is now 0644
1618 '''
1619 The old 'mode=000' inode is unlinked by 'mv'; the new 'mode=0644' inode (created by base64-decode via the
    calling user's umask) takes the path. No chmod ever called. The parent directory just needs to be
    writable by you (it always is, since you own it).
1620
1621 **Flavor B locked file is mode '0444' (you can READ but cannot WRITE; 'scripts/sync-rules.py' needs to
    re-write):**
1622 You own the file and can read it, you just cannot truncate/overwrite it. The simplest path is to rename
    it aside and let the writer create a fresh inode from scratch (which gets the umask-default '0644'):
1623 '''bash
1624 mv <FILE> /tmp/old.<FILE>.$$
1625 python3 scripts/sync-rules.py # now writes <FILE> cleanly at mode 0644

```

```

1626 '''
1627 Or apply Flavor A if you'd prefer to preserve the contents while resetting the mode.
1628
1629 **Why this is the canonical recovery, not a hack:**
1630 - Mode bits live in the inode, not the path. 'mv -f' replaces the path's inode reference; the old locked
    inode drops out of the directory entirely (dentry eviction) and loses its last link, so the kernel
    reclaims it. The new inode (whatever it is: a freshly-written one, or a base64-decoded copy) gets
    the path.
1631 - The only prerequisite to apply Flavor B's 'mv' is: you own the parent directory (so you can unlink
    entries from it) AND you have a NOPASSWD-capable binary that can read the byte stream if mode
    prohibits user read.
1632 - This pattern generalizes. Any time a script ends up producing a "locked file that the next run can't
    update" situation, the same trick applies without a single chmod anywhere.
1633
1634 **Empirical verification (user machine, July 1, 2026):**
1635 - '.env' was 'mode=000' (owner $USER, i.e. you): Flavor A unblocked it in 3 commands, no chmod.
1636 - Before: '----- 1 $USER staff 899 Jun 28 07:07 .env'
1637 - After: '-rw-r--r--@ 1 $USER staff 899 Jul 1 20:39 .env'
1638 - 'docker compose config' immediately picked up 'TS_AUTHKEY' + 'WIREGUARD_PRIVATE_KEY' substitution.
1639 - 'compiled_rules.txt', 'ip_blocklist.ipset', 'data/filters/1782645604.txt' were 'mode=0444': Flavor B ('
    mv'-aside) unblocked all three.
1640 - scripts/sync-rules.py then ran clean: 279587 block rules + 171 allow rules + 16710 IP entries
    compiled; new files came out at '0644'.
1641 - 'toggle.sh' then went end-to-end in 48s (warp / routing-fix / adguardhome / tailscale all healthy,
    gateway mesh-joined, DNS hijack verified single-server).
1642
1643 **When this trick is appropriate vs. inappropriate:**
1644 - You own the parent directory and the file (typical desktop scenario).
1645 - The lock is a leftover from a removed 'chmod' call that you want off disk forever.
1646 - NOPASSWD sudo includes at least one interpreter ('python3' on this system) that can be used to read
    mode=0 files. If it does not, add it first: '<USER> ALL=(ALL) NOPASSWD: /usr/bin/python3'.
1647 - The file is owned by another user (root, another account) and the lock is intentional respect it; do
    not bypass another user's security boundary.
1648 - The parent directory is owned/writable only by another user escalate properly via sudo, or stop and
    ask the user.
1649 - You do not actually own the file and there is a legitimate reason for its lock state restore the file
    via a trusted source rather than circumventing the perms.
1650
1651 **Reusable cookbook:**
1652
1653 '''bash
1654 # Quick diagnostic: figure out which flavor you need.
1655 WHOAMI=$(whoami)
1656 ls -la <FILE>
1657 stat -f 'mode=%Sp owner=%Su group=%Sg' <FILE>
1658 # mode=----- or mode=---r----- : you cannot even read it Flavor A.
1659 # mode=r--r--r-- but writer fails : you are blocked from write but can read Flavor B.
1660 # mode=rw-r--r-- : not actually locked at all; unrelated write failure.
1661
1662 # Flavor A (mode=000):
1663 sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb')).
    read()).decode())" > /tmp/snap.b64
1664 base64 -d < /tmp/snap.b64 > <FILE>.new && mv -f <FILE>.new <FILE> && rm -f /tmp/snap.b64
1665
1666 # Flavor B (mode=0444, file is readable):
1667 mv <FILE> /tmp/old.<FILE>.$(date +%s)
1668 # (let the original writer recreate <FILE>; new file lands at mode 0644)
1669
1670 # Verify after either flavor:
1671 ls -la <FILE> # mode should be 0644
1672 <your normal validation command>
1673 '''
1674
1675 ### 15.10 Tor Container
1676
1677 #### 15.10.1 July 11, 2026: Tor Container Hardening & obfs4proxy Restoration
1678 **Symptom:**
1679 1. The Tor container entered a fatal crash loop upon boot, throwing 'exec: "obfs4proxy": executable file
    not found in $PATH' when 'TOR_USE_BRIDGES=true'.
1680 2. When the user pasted Tor bridge lines containing comments ('#') or blank lines into 'tor-bridges.txt',
    the container failed to validate the file properly and crashed.
1681
1682 **Root Cause:**
1683 - **Bug 1 (Missing Pluggable Transports):** The Tor Dockerfile was built on 'debian:bullseye-slim', which
    had silently dropped or failed to resolve the 'obfs4proxy' package in its latest apt repositories.
1684 - **Bug 2 (Fragile Validation):** The 'entrypoint.sh' bridge validation check merely checked '[ -s tor-
    bridges.txt ]' (file size > 0). It did not verify if the file actually contained valid, uncommented
    bridge lines. Passing empty lines or comments tricked the validation script into appending invalid '
    Bridge' directives to the 'torrc', causing the Tor daemon to instantly exit.
1685
1686 **Resolution:**

```

```

1687 - **Upgraded Base Image:** Shifted the Tor Dockerfile base image to 'debian:bookworm-slim', restoring
      access to the 'obfs4proxy' pluggable transport package.
1688 - **Robust Bridge Validation:** Replaced the fragile '[ -s ]' check with a strict 'grep -vE '\s*(#|$\)'
      command that strips comments and empty lines. If no valid lines remain, the container safely falls
      back to standard public guard relays instead of crashing, ensuring Tor remains available even with a
      misconfigured bridge file.
1689 - **Image Pinning:** Explicitly pinned 'image: nullexit-tor:v1.0.0' in 'docker-compose.yml' to prevent
      arbitrary local builds from drifting to ':latest'.
1690
1691 ##### 15.10.2 July 12, 2026: Tor ControlPort Dynamic Password Authentication & User Config
1692 **Symptom:** Querying the Tor Control Port (e.g., via the '/sweep' script or 'check_circuits.py')
      returned empty responses or connection errors. Logs showed that Tor automatically closed its
      ControlPort:
1693 'You have a ControlPort set to accept unauthenticated connections from a non-local address. ... That's so
      bad that I'm closing your ControlPort for you.'
1694
1695 **Root Cause:** Docker port mapping (especially under Colima on macOS) requires container sockets to bind
      to '0.0.0.0' to be reachable from the host. However, when Tor's ControlPort is bound to '0.0.0.0' (
      non-local) with no authentication, Tor closes it by default as a safety precaution. Sourcing SOCKS5/
      ControlPort to '127.0.0.1' inside the container was not an option as it broke the Docker bridge
      routing path.
1696
1697 **Resolution:**
1698 - **Dynamic Password Authentication:** Updated the Tor container's 'entrypoint.sh' to read 'TOR_PASSWORD'
      and dynamically generate its HMAC hash via 'tor --hash-password', adding 'HashedControlPassword <
      hash>' to 'torrc'. This secures the ControlPort on '0.0.0.0' so Tor keeps it open.
1699 - **Unified Credentials in '.env':** Exposed 'ADGUARD_USER=admin' and 'ADGUARD_PASSWORD=nullexit' in '.
      env' to serve as the unified gateway credentials. Sourced 'TOR_PASSWORD' directly from '
      ADGUARD_PASSWORD' in 'docker-compose.yml'.
1700 - **Client Authentication:** Updated 'scripts/check_circuits.py' to fetch 'TOR_PASSWORD' (defaulting to '
      ADGUARD_PASSWORD' or 'nullexit') and authenticate with the ControlPort using 'AUTHENTICATE "<
      password>"' to regain access.
1701
1702 > See also 11.3 (Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification) for the design-
      level fail-closed guarantee.
1703
1704 ##### 15.11 macOS Platform Quirks
1705
1706 ##### 15.11.1 macOS Application Firewall Silently Blocking AirDrop & Continuity
1707 **Problem:** Users of complex network extensions (Tailscale, LuLu, Docker, etc.) often run into an issue
      where AirDrop, AirPlay, and Universal Clipboard silently fail, even with the gateway inactive and Wi
      -Fi/Bluetooth correctly configured. The internal Apple Wireless Direct Link ('awdl0') interface
      appears healthy, but connections never establish.
1708
1709 **Root Cause:** The built-in macOS Application Firewall has a specific list of explicitly blocked
      applications. Apple's own core networking services (e.g., '/usr/libexec/sharingd', '/usr/libexec/
      rapportd', '/usr/libexec/avconferenced') can be silently added to the "Block incoming connections"
      list either through accidental "Deny" clicks on permission prompts, execution of aggressive privacy-
      hardening scripts, or a known macOS bug that retains strict blocks even after toggling the "Block
      all incoming connections" setting off. Since these are background daemons, macOS provides no visible
      error.
1710
1711 **Fix:** The firewall rules must be cleared via the terminal (or System Settings > Network > Firewall >
      Options):
1712 ```bash
1713 sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/sharingd
1714 sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/rapportd
1715 ```
1716 Once unblocked, 'sharingd' immediately resumes broadcasting mDNS/Bonjour discovery packets, and AirDrop
      restores instantly without requiring a reboot.
1717
1718 ##### 15.11.2 Standalone Tailscale Daemon Freeze after macOS Sleep/Wake
1719 **Problem:** When the macOS host goes to sleep and wakes up, the network interfaces and routing tables
      are rebuilt. The standalone 'tailscaled' daemon (installed via Homebrew) occasionally fails to
      handle this transition and becomes completely frozen/unresponsive. In this state, any command like '
      tailscale down' or 'tailscale status' hangs indefinitely, and all DNS requests sent to the local
      Tailscale resolver ('100.100.21.8') time out, causing a total internet blackout on the host.
1720
1721 **Fix:** Enhanced the Tailscale verification and teardown logic in 'toggle.sh' and 'recover.sh':
1722 1. **Unresponsive Daemon Detection:** Instead of assuming the daemon is healthy just because the Unix
      socket '/var/run/tailscaled.socket' exists, the scripts now actively run a status check with a 5-
      second timeout ('run_with_timeout 5 tailscale status'). If the check times out (exit code 143), the
      daemon is classified as unresponsive.
1723 2. **Auto-Recovery Loop:** If the daemon is unresponsive, the script now automatically attempts to
      restart the service using 'brew services restart tailscale' (falling back to 'sudo -n brew services
      restart tailscale').
1724 3. **Graceful Retry:** Once restarted, the script pauses for 3 seconds for the daemon to initialize
      before proceeding with connection or teardown commands.
1725
1726 ##### 15.11.3 macOS Sharing Services (AirDrop/AirPlay) Freeze on Network Transitions

```

```

1727 **Problem:** When the gateway is turned on or off, the scripts perform network configuration cleanups
      which include flushing the routing tables ('route -n flush'), restarting the 'mDNSResponder' daemon
      ('killall -HUP mDNSResponder'), and power-cycling the Wi-Fi interface ('setairportpower' or '
1728 ifconfig en0 down/up'). While AirDrop BLE discovery continues to function, the actual file transfers
1729 stall at "Waiting..." indefinitely.

**Root Cause:** The rapid tear-down and rebuild of the local routing table and Wi-Fi interface causes
      Apple's core sharing and connection daemons ('sharingd' and 'rapportd') to lose their socket
      bindings to the mDNS/Bonjour discovery interface. Instead of reconnecting gracefully, they enter a
      wedged state and try to route peer-to-peer (AWDL) IPv6/TCP transfer traffic into the default gateway
      (the Tailscale interface 'utun') instead of the direct wireless link, causing the TLS verification
      handshake to time out.

1730 **Fix:** Integrated an automatic sharing services reset into both 'toggle.sh' and 'recover.sh':
1731 1. The script now automatically runs 'sudo -n killall sharingd rapportd' at the end of both the **START**
1732 path and the **STOP** path (as well as inside 'recover.sh').
1733 2. Killing these daemons forces macOS to immediately relaunch them, binding them fresh to the newly
      initialized network interfaces and routing tables, allowing AirDrop to work seamlessly while the
      gateway is active.

1734 #### 15.11.4 Chrome Remote Desktop Connection Failures (Tailscale on Remote Device)
1735 **Context:** When attempting to access a remote machine via Chrome Remote Desktop (CRD), the connection
1736 fails or hangs if Tailscale is active **on the remote device**. The connection works fine even if
      Tailscale (and the exit node) is active **on the local client/viewer device**, but only if Tailscale
      is disabled on the remote machine. *(This is distinct from the CRD-through-WARP latency limitation
      in 14.3.)*

1737 **Why:** Having Tailscale active on the remote machine creates virtual network interfaces and DNS/routing
1738 modifications that conflict with Chrome Remote Desktop's WebRTC bindings and signaling traffic.

1739 **Workaround:**
1740 - **Use Tailscale Native RDP/RustDesk (Recommended):** Connect directly to the remote device's Tailscale
1741 IP ('100.X.Y.Z') via an RDP or RustDesk client. This is faster and avoids third-party relay
      conflicts.
1742 - **Disable Tailscale on the Remote Device:** If you must use CRD, disable Tailscale on the remote
      machine before connecting.

1743 #### 15.11.5 tailscaled Daemon Broken State (brew services)
1744 **Symptom:** 'brew services list' shows 'started' but 'tailscale status' shows "stopped".
1745 **Fix:** 'sudo brew services restart tailscale'. Toggle script now detects and auto-restarts.

1746 #### 15.11.6 Apple Silently Removed the 'airport' Binary in macOS Sonoma 14.4 (March 2024)
1747 **What happened:** 'watcher.sh's 'detect_lan_p2p_mode()' was written to use the 'airport' CLI tool at:
1748 '''
1749 /System/Library/PrivateFrameworks/Apple80211.framework/Versions/Current/Resources/airport
1750 '''
1751 This path returned 'exit 127: no such file or directory'. The function silently fell back to 'reason="no-
1752 wifi"' even while the Mac was actively connected to WPA2 Enterprise Wi-Fi.

1753 **Why Apple removed it:** 'airport' was never a real public binary it lived inside a **private framework
1754 ** and was always technically unsupported. Apple deprecated it quietly years ago and finally removed
1755 it in **macOS Sonoma 14.4 (March 2024)**. The removal is part of a broader privacy clampdown on Wi-
1756 Fi metadata:
1757 - Starting in **macOS 14.5**, BSSIDs and other Wi-Fi association details are **redacted** ('<redacted>')
1758 in most tools, including the official replacement 'wdutil'
1759 - Apple's official recommendation is to use the **CoreWLAN framework** (Swift/Objective-C) with proper
      entitlements useless for a bash script
1760 - The official CLI replacement 'wdutil' requires 'sudo' for most useful output also useless for a
      background LaunchAgent

1761 **Fix:** Replaced 'airport -I | awk '/link auth/' with:
1762 '''bash
1763 system_profiler SPAirPortDataType | awk '/Current Network Information:/{found=1} found && /Security:/{
1764 print $NF; exit}'''
1765 This returns human-readable strings like 'WPA2 Enterprise', 'WPA2 Personal', or 'None' for the currently
      associated network without 'sudo', and without any removed private binary.

1766 **Confirmed working output on this machine:**
1767 '''
1768 P2P detect: security='Enterprise' allow=false (reason: 802.1x-enterprise)
1769 '''
1770 > ** Risk: 'system_profiler' may not survive forever either. ** 'system_profiler SPAirPortDataType' still
1771 works as of macOS Sequoia (15.x) without sudo and without redaction. However, given Apple's
1772 trajectory on Wi-Fi privacy, this could be locked down in a future release. If 'security' comes back
      empty on a future macOS version while the Mac is actively on Wi-Fi, the function will silently fall
      back to 'allow=false' (safe default), but the detection will be broken again. The long-term fix
      would be a small Swift helper using CoreWLAN that we compile once and bundle in the repo. *(This
      forward-looking caveat is also tracked as an observation in 16.)*

```

```

1773 ##### 15.11.7 Changing macOS Local Hostname
1774 **Context:** Users may wish to change their Mac's computer name/hostname in macOS System Settings.
1775 **Effect:** Changing the Mac's hostname will not break Nullexit (which relies on internal routing/
1776 localhost). However, Tailscale will automatically detect this change and update the machine name on
the Tailnet.
1777 - The underlying Tailscale '100.x.x.x' IPv4 address will remain exactly the same.
1778 - The MagicDNS name will change to match the new hostname. Users must remember to update their SFTP/
SSH clients to use the new MagicDNS address.
1779 ##### 15.11.8 macOS MAC Address Spoofing (Apple Silicon)
1780 **Observation:** macOS prevents setting a new MAC address ('ifconfig en0 ether <mac>') while the Wi-Fi
1781 interface is actively associated with an SSID. Additionally, some drivers on Apple Silicon
completely reject spoofed MAC addresses via 'ifconfig', effectively ignoring the command and
retaining the hardware MAC.
1782 **Workaround / Fix:**
1783 1. Disassociate First: The Wi-Fi interface must be turned off ('networksetup -setairportpower en0 off
') before attempting to set the MAC address, and turned back on afterward. This allows the OS to
accept the change.
1784 2. Fail-Safe Testing: The 'scripts/device-scramble.sh test-mac' command was implemented to safely
verify if the underlying hardware and the current network will accept a randomized MAC address, as
behavior varies significantly between Intel Macs (usually permissive) and Apple Silicon Macs (often
restrictive).
1785 ##### 15.11.9 Shell-Language Landmines: macOS '/bin/bash' 3.2 + 'set -e' (Field Notes)
1786 This project's lifecycle scripts ('toggle.sh', 'recover.sh', 'common.sh', ) run under '#!/bin/bash',
1787 which on macOS is GNU bash 3.2.57 (2007). Apple has never shipped a newer bash because bash 4+
is GPLv3. (Homebrew's bash 5 lives at '/opt/homebrew/bin/bash' but is not the shebang target, and
cannot be assumed on PATH - a bare 'bash --version' on a stock Mac reports 3.2.) Every script here
also runs 'set -e' (exit-on-error) with an 'ERR'/'INT'/'TERM'/'HUP' trap. That combination ancient
bash + 'set -e' + traps + 'set -x' xtrace redirection is a minefield. The traps below
were all hit and fixed during development; document them so they are not re-discovered.
1789 A. 'set -e' + 'if then else fi' can abort at the 'else' (bash 3.2 heisenbug).
1790 The nastiest one. A construct as innocent as:
1791 ''bash
1792 if [ -f "$MARKER_FILE" ]; then
1793     X="yes"
1794 else
1795     X="no"
1796 fi
1797 ''
1798 aborted 'toggle.sh' at init under 'set -e' whenever the file was absent the false status of the '[ -
f ]' test leaked through the compound and 'set -e' killed the script ('EXIT: code=1 phase=init'),
before any real work ran. Two things made this vicious: (1) it is intermittent/context-dependent it
did not reproduce in isolated 'bash -c '' snippets of the identical block; it only manifested
inside the full script with the 'set -x' xtrace fd-redirect ('exec 2> >(tee )') active. We only
proved causation by bisecting the live script (replace the block with a trivial assignment the
script sailed past init). (2) Static tools ('bash -n', 'shellcheck') cannot see it it is a runtime
interaction. Safe pattern: default-assign, then an 'if' with no 'else' (an 'if' whose false
condition has no else-branch yields status 0):
1800 ''bash
1801 X="no"
1802 if [ -f "$MARKER_FILE" ]; then X="yes"; fi
1803 ''
1804 Confirmed root cause (2026-07-13, reproduced head-to-head). The trigger is specifically 'set -x'
plus a 'PS4' that contains a command substitution here the 'DEBUG_TRACE' prompt 'PS4='+ [$(date) ]
'. Forcing the false-condition/else path ('gateway_state' 'stopped', so 'if [ "$(gateway_state)"
= "running" ]' takes the else/START branch) on stock '/bin/bash' 3.2.57: with the default 'PS4'
it survives (3/3, exit 0); with toggle's '$(date)' 'PS4' the 'ERR' trap fires on the else
branch ('rc=1' from the false condition, 3/3, exit 42). The '$(date)' runs a subshell on every
traced command; while tracing the 'if'-condition it perturbs '$?' so the condition's non-zero status
leaks past the normal 'if'-exemption and trips the 'ERR' trap. This is a true heisenbug
enabling tracing ('DEBUG_TRACE=true') is what causes the abort, which is exactly why it's
invisible with 'DEBUG_TRACE=false' (the default) and why every isolated repro on the default 'PS4'
passes (this misled an entire debugging session into briefly "disproving" the landmine). It aborted
the gateway START decision ('toggle.sh:689') under 'DEBUG_TRACE=true'; with 'DEBUG_TRACE=false' the
gateway starts cleanly (verified: 192 s, all five containers healthy). Fixed (2026-07-14).
DEBUG_TRACE's 'PS4' now uses the subshell-free '${SECONDS}' (elapsed seconds since the shell
started) instead of '$(date +%H:%M:%S)', removing the per-traced-command subshell that perturbed
'$. bash 3.2 has no subshell-free wall-clock token for 'PS4' ('$EPOCHSECONDS' is 5.0+, 'printf
'$(T)' is 4.2+), so the per-line stamp is now elapsed-seconds the absolute start time is still
logged once in the 'DEBUG_TRACE enabled' header line, so wall-clock is recoverable. Bonus: this
also drops a 'date' fork on every traced line (previously thousands). Verified head-to-head on
stock '/bin/bash' 3.2.57: 'toggle.sh --restart' with 'DEBUG_TRACE=true' now completes 'code=0'
with tracing active straight through the exact decision that aborted before the trace shows
gateway_state stopped ('common.sh:229') feeding the false 'if [ "$(gateway_state)" = "running" ]',
then the else branch reaching Action: STARTUP GATEWAY (1436 traced lines, no 'ERR' trap firing,
all five containers healthy). 'DEBUG_TRACE' is now safe to use on the STOP/START decision path.

```


1806 More generally, any command whose "failure" is normal control flow ('grep -q', '[]', 'diff', 'kill -0')
that ends up as the *last* command of a function/branch/script under 'set -e' is a landmine; guard
it ('|| true') or restructure so a non-zero status can't be the tail value.

1807

1808 ****B.** 'BASH_XTRACEFD' and the 'exec {var}>>file' dynamic-fd form are bash 4.1+ (hard-fail on 3.2).******
1809 Our 'DEBUG_TRACE' feature wants xtrace ('set -x') sent to 'output.log' while keeping the terminal clean
the clean way is 'exec {BASH_XTRACEFD}>>"\$LOG_FILE"'. On 3.2 this is a ****syntax/runtime error****:
'exec: {BASH_XTRACEFD}: not found' (dynamic-fd allocation '{var}>>' didn't exist until 4.1, and
'BASH_XTRACEFD' itself is 4.1+; on 3.2 it's just an ordinary variable that 'set -x' ignores). Left
unguarded it would abort the script at startup i.e. the debug feature would break the very thing it
's meant to observe. ****Safe pattern:**** branch on '{BASH_VERSINFO[0]}.\${BASH_VERSINFO[1]}' on 4.1
use 'exec 9>>"\$LOG_FILE"; export BASH_XTRACEFD=9'; on 3.2 fall back to 'exec 2> >(tee -a "\$LOG_FILE"
>&2)' (xtrace goes to stderr, which we duplicate to the log; the terminal also shows it, which is
acceptable). Never use the 'exec {var}>>' form unless you ***know*** you're on 4.1.

1810

1811 ****C.** 'trap ERR' + '\$LINENO' misattributes the failing line on 3.2.******
1812 When a command fails under 'set -e', our trap runs 'cleanup_handler ERR \$LINENO'. On bash 3.2, '\$LINENO'
inside an 'ERR' trap frequently reports the line of the ****enclosing 'fi'/closing brace of a compound**
statement**, not the true failing line. During the STOP/START-decision bug this produced 'Script
failed at line 1202' where 1202 is a bare 'fi' actively misleading. ****Mitigation:**** don't trust the
trap's line number as gospel on macOS; corroborate with the 'EXEC:/EXIT:' lifecycle breadcrumbs
and the 'DEBUG_TRACE' xtrace (which prints 'file:line:function' via a custom 'PS4'), which pin the ***
actual* last-executed command**.

1813

1814 ****D.** 'set -e' is not inherited by command substitutions / subshells the way people expect, and is
suppressed inside 'if/`&&`/||' conditions.******
1815 This cuts both ways and is a frequent source of "why didn't it fail?" and "why ***did*** it fail?": a non-
zero command used as an 'if' condition (or left/right of '&&`/||') is ***exempt*** from 'set -e' (by
design), but the moment that same call is a bare statement it aborts. 'is_gateway_active()' relies
on this (it's always called as 'if is_gateway_active; then '), which is precisely why moving the
decision to a marker file instead of a probe that must be 'if'-guarded to be safe is more robust (see
15.12.19). Also note 'local X=\$(cmd)' masks 'cmd's exit status (the 'local' builtin's status
wins) a 'shellcheck' SC2155 we see repeatedly; harmless when we don't check '\$?', but a real
footgun if you do.

1816

1817 ****E.** Process substitution ('>()','<()') is a bashism, and its lifecycle is loose.******
1818 Our 3.2 xtrace fallback ('exec 2> >(tee -a "\$LOG_FILE" >&2)') uses process substitution fine under bash,
but (1) it is ****not**** POSIX 'sh', so these scripts must never be invoked as 'sh script.sh'; and (2)
the background 'tee' it spawns is reaped asynchronously, so the very last few lines written just
before exit can occasionally be truncated in the log. Acceptable for a debug trace; do not rely on
it for critical last-gasp logging (use a direct '>> "\$LOG_FILE"' append for must-capture lines, as
the 'EXIT' breadcrumb does).

1819

1820 ****F.** GNU-vs-BSD userland divergence (not bash itself, but same class of pain).******
1821 macOS ships BSD variants of the core tools, so idioms differ from Linux: 'sed -i '' (BSD requires the
empty backup-suffix arg) vs 'sed -i' (GNU); 'stat -f%z' (BSD) vs 'stat -c%s' (GNU); 'date' format
flags; 'jot' (BSD) vs 'shuf' (GNU) for random port generation; 'route'/'networksetup'/'dscacheutil'
(macOS-only) vs 'ip'/'resolvectl' (Linux). This is ***why*** the codebase branches on '[["\$OSTYPE" == "
darwin"*]]' everywhere 'toggle.sh' and 'recover.sh' are each single cross-platform scripts that
dispatch on '\$OSTYPE' (an early Linux block runs and exits before the macOS body). 'timeout(1)'
doesn't exist on stock macOS at all hence the pure-bash 'run_with_timeout' in 'common.sh'.

1822

1823 ****Working rules distilled from the above:****
1824 - Assume ****bash 3.2**** semantics for anything under '#!/bin/bash' on macOS; gate any 4.0+ feature ('
BASH_XTRACEFD', dynamic '{var}' fds, associative arrays, '\${var^^}' case-conversion, 'mapfile'/'
readarray', negative array indices) behind a 'BASH_VERSINFO' check or avoid it.
1825 - Under 'set -e', never let a "benign non-zero" command ('[]', 'grep -q', 'kill -0', 'diff') be the
tail statement of a branch/function/script; use no-'else' 'if's, '|| true', or explicit '\$?'
handling.
1826 - Prefer ****persisted state (marker files)**** over ****live probes**** for control-flow decisions a probe
forces an 'if'-guard and can hang/abort; a file read cannot (see 15.12.19).
1827 - 'bash -n' and 'shellcheck' catch syntax and many smells but ****cannot**** catch 'set -e'/trap runtime
interactions a ****live run is mandatory**** to validate lifecycle-script changes; the 'DEBUG_TRACE' +
breadcrumb instrumentation exists precisely to make those live runs diagnosable.
1828 - Never run these scripts as 'sh' they are bashisms end to end.

1829

1830 **### 15.12 Misc Resolved**

1831

1832 **#### 15.12.1 Gluetun Resets nftables FORWARD Rules**
1833 ****Symptom:**** FORWARD RELATED, ESTABLISHED rule disappears after Gluetun health check.
1834

1835 ****Fix:**** 'scripts/routing-fix.sh' re-adds to both iptables backends every 30 seconds. Legacy backend (policy ACCEPT) acts as safety net.

1836

1837 **#### 15.12.2 Native SSH & SFTP Security (Zero Local Attack Surface)**
1838 ****Context:**** macOS "Remote Login" ('sshd') and "File Sharing" ('smbd') open ports on the local network (e .g. '172.x.x.x'), exposing the host to brute-force attacks on public Wi-Fi.
1839

1840 ****Implementation:**** The script strictly enforces 'tailscale up --ssh=true' whenever the mesh state is
reset. This empowers the user to completely disable native macOS Remote Login and File Sharing. This
provides an incredible zero-trust security advantage: even if an attacker steals your Mac username

and password, they **cannot** access your files remotely. Disabling the native services completely closes the listening ports on the local Wi-Fi interface ('172.x.x.x'), rendering the Mac inaccessible to local network logins. Meanwhile, Tailscale intercepts port 22 traffic exclusively over the encrypted mesh and authenticates cryptographically based on your Tailscale SS0 identity. Stolen Mac passwords are mathematically useless without first bypassing your Tailscale Two-Factor Authentication and registering a device onto the mesh.

****Cross-Platform File Exchange:**** To easily browse and exchange files securely, simply use an SFTP client and connect to your Mac's MagicDNS name on port 22. Recommended clients: ****WinSCP**** or ****FileZilla**** (Windows), ****Cyberduck**** (macOS), and ****FE File Explorer**** or ****Solid Explorer**** (iOS/Android).

15.12.3 Logging Architecture

For debugging, logs are strictly segmented based on the component's lifecycle:

- ****'output.log' (Host-side):**** Contains all standard error ('stderr') and verbose output from the host scripts ('toggle.sh', 'recover.sh', 'setup.sh'). Since the 'rule-compiler' container is ephemeral and deleted after running, its logs (and any Python errors from 'scripts/sync-rules.py') are extracted and appended to this file before deletion.
- ****Log Rotation Policy:**** On startup, 'toggle.sh' checks if 'output.log' exceeds ****1GB**** ('1,073,741,824' bytes). If it does, it performs a metadata-only rename ('mv output.log output.log.old'), preserving history while resetting the active log to 0 bytes. This rename-based rotation avoids the heavy disk I/O and CPU overhead of rewriting a massive file line-by-line. The threshold is intentionally large because with 'DEBUG_TRACE=true' the log accumulates a full command-by-command execution trace (effectively the framework's entire compilation/warning history), which is valuable to retain for post-mortems; a smaller cap would rotate that history away too quickly.
- ****Egress Probe Logging:**** The background host-egress leak probe checks if stdout is a TTY ('[-t 1]') and will only print live status updates (using carriage return '\r') or duplicate console warnings in interactive mode. It automatically detects macOS/BSD vs. Linux environments to dynamically construct timestamps with date and time (omitting milliseconds on macOS to prevent high CPU utilization from spawning sub-second processes).
- ****'docker logs <container>' (Guest-side):**** The persistent containers ('warp', 'tailscale', 'routing-fix', 'adguardhome') use the Docker 'json-file' logging driver with a strict 'max-size' (1m-10m) to prevent VM disk exhaustion. Use standard 'docker logs' to view them.

15.12.4 Pre-Flight Check 1: Wrong 'tailscale ping' Flag

****Problem:**** Check 1 of the pre-flight connectivity checks used '--c 1' (double dash) but 'tailscale ping' accepts '-c 1' (single dash). The invalid flag caused the command to exit with code 1, marking the check as FAIL even though the gateway was reachable.

****Fix:**** Changed '--c 1' to '-c 1'.

15.12.5 Pre-Flight Check 1: DERP Relay Exit Code

****Problem:**** Even with the correct '-c 1' flag, 'tailscale ping' exits with code 1 when the connection goes through a DERP relay ('direct connection not established') even though a pong was successfully received (28ms via DERP(tor)). The check treated relayed connections as failures.

****Fix:**** Changed the check from relying on exit code to piping stdout through 'grep -q "pong"'. This accepts relayed connections as valid (they work fine for the exit node use case), while still failing on true timeouts or unreachable peers.

15.12.6 IPv6 Exit-Node Leak

****Problem:**** Tailscale supports IPv6, but WARP/gluetun is IPv4-only. IPv6 packets forwarded through the exit node would leak through the Docker bridge unencrypted.

****Fix:**** Added 'ip6tables -A FORWARD -i tailscale0 -j DROP' to 'post-rules.txt', blocking all IPv6 forwarded traffic from the Tailscale interface. *(This is the fix referenced by the IPv6-leak scenario in 15.6.1.)*

15.12.7 SOCKS5 Proxy Threading Race Condition

****Problem:**** The 'forward' function in 'scripts/socks5-proxy.py' called 'select.select([src, dst], [], [])' which raised 'ValueError: file descriptor cannot be a negative integer (-1)' when one thread closed a socket while the other thread was selecting on it.

****Fix:**** Added 'ValueError' exception handling around 'select.select()' and 'recv()' calls. When a file descriptor becomes negative, the thread breaks out of the forwarding loop cleanly instead of crashing.

15.12.8 Docker Healthcheck: Multi-line Python in CMD Array

****Problem:**** The socks-proxy healthcheck used a 'CMD' array with a multi-line Python string. YAML collapsed the newlines, causing the '#' comment to comment out the rest of the code and the 'assert' statement to be mangled into 'as'.

****Fix:**** Changed from '["CMD", "python3", "-c", "..."]' to '["CMD-SHELL", "python3 -c '...'"]' with a single-line Python expression. Uses a real SOCKS5 handshake (sends '[5, 1, 0]', expects '[5, 0]') instead of a simple TCP port check.

15.12.9 Graceful Shutdowns Leave DNS Hijacked

****Problem:**** The gateway overrides the macOS host's DNS settings (via 'networksetup') to route queries to the AdGuard container. Because macOS persists 'networksetup' changes to disk, shutting down the Mac while the gateway is running causes the Mac to boot up later with hijacked DNS. Since the Colima VM and Docker containers don't auto-start on boot, the user would have no internet access until they manually run the toggle script.

1878

1879 ****Root Cause:**** The 'toggle.sh' script exits after the gateway starts. There was no background daemon listening for macOS system shutdown signals to revert the network settings.

1880

1881 ****Fix:**** Wrapped the background 'caffeinate' process (used to prevent system sleep) in a bash 'trap' command. The trap listens for 'SIGTERM', 'SIGINT', and 'SIGHUP'. When the user initiates a graceful shutdown, macOS sends 'SIGTERM' to the background trap, which instantly executes 'recover.sh' to flush the host DNS back to '1.1.1.1' right before the machine powers off.

1882 ***Note:*** We explicitly use 'recover.sh' instead of 'toggle.sh' because during a shutdown, the OS limits process cleanup time and is already independently sending termination signals to Docker and Colima. Trying to run 'docker compose down' inside a shutdown trap creates a race condition that can hang the script until macOS forcefully 'SIGKILL's it. 'recover.sh' abandons container management and surgically flushes the macOS network settings in milliseconds, guaranteeing completion before the OS pulls the plug.

1883

1884 **#### 15.12.10 Linux Native Wi-Fi Bouncing**

1885 ****Problem:**** The generated 'scripts/recover-linux.sh' incorrectly retained the macOS 'networksetup' commands, which would crash and fail to bounce Wi-Fi on Linux hosts.

1886

1887 ****Fix:**** Swapped the macOS logic for native Linux commands. The script now attempts 'nmcli radio wifi off /on' first, and gracefully falls back to 'ip link set wl... down/up' if NetworkManager is unavailable.

1888

1889 **#### 15.12.11 TS_AUTH_ONCE Cloud Footgun & Ephemeral Keys**

1890 ****Decision:**** The compose file uses 'TS_AUTH_ONCE=true' to prevent authentication loops. However, on headless cloud VPS deployments, if a standard auth key expires (usually 90 days), the container will silently drop off the mesh. We documented this trade-off explicitly in 'docker-compose.yml' and recommended the use of ****Tailscale Ephemeral Keys**** for cloud deployments, which automatically clean themselves up and bypass this issue.

1891

1892 **#### 15.12.12 Hardcoded AdGuard Credentials**

1893 ****Decision:**** AdGuard is deployed with the hardcoded credentials 'admin / nullexit'. This is perfectly safe behind a NAT on a local macOS laptop. However, if deployed on a cloud host with exposed ports, this becomes a severe security risk. We added a stark warning comment in 'docker-compose.yml' to ensure cloud users change this before deployment.

1894

1895 **#### 15.12.13 Routing Fix: 30-second polling vs Netlink Socket**

1896 ****Decision:**** Instead of building a complex Netlink socket listener ('ip monitor') to react instantly to iptables and routing flushes from Gluetun, we retained the 5-second 'sleep' loop in 'scripts/routing-fix.sh'.

1897 **- **Reasoning:**** Building a netlink listener in pure bash on Alpine is incredibly brittle and complex (especially for iptables events). The 30-second polling loop is bulletproof. The only trade-off is a potential 1-4 second stutter if Gluetun reconnects, which was deemed an acceptable edge case for absolute reliability.

1898

1899 **#### 15.12.14 Cache Poisoning and URL Rot in scripts/sync-rules.py**

1900 ****Problem:**** If a remote blacklist URL went permanently offline (404 Not Found), the Python 'urllib' request would return the 404 HTML string. The script would aggressively regex-parse this HTML, find no domains, and overwrite the healthy local cache with a 0-domain file. This effectively disabled ad-blocking until the URL was fixed.

1901

1902 ****Fix:**** Implemented a defensive programming sanity check in 'scripts/sync-rules.py'. Before overwriting the cache, the script now verifies that the compiled domain count is greater than 10 (and 1 for IP feeds). If it falls below this threshold (indicating a catastrophic 404 failure across multiple URLs), it aborts the cache overwrite, raises a 'ValueError', and keeps the previous day's healthy cache intact.

1903

1904 **#### 15.12.15 Gluetun nf_tables Parser Crash (Bash Interpolation Failure)**

1905 ****Problem:**** Attempting to make the TCP MSS dynamically configurable by using a standard bash variable inside 'post-rules.txt' ('iptables ... --set-mss \${GATEWAY_MSS:-1120}') caused the Gluetun container to catastrophically crash during startup. Gluetun's internal parser executes 'post-rules.txt' line by line without a shell, meaning the variable was never expanded. It literally attempted to inject the string '\${GATEWAY_MSS:-1120}' into iptables, resulting in a fatal 'bad value for option "--set-mss"' error.

1906

1907 ****Fix:**** Removed the bash variable from 'post-rules.txt' and reverted it to a pure hardcoded integer. To retain dynamic '.env' configuration, we moved the interpolation logic to 'toggle.sh'. Right before Docker starts, the host script parses 'GATEWAY_MSS' from the '.env' file and uses a cross-platform 'sed' command to dynamically overwrite the integer directly inside 'post-rules.txt'. This perfectly mimics manual configuration and completely bypasses Gluetun's strict parser.

1908

1909 **#### 15.12.16 AdGuard Filter Redundancy & Memory Optimization**

1910 ****Problem:**** AdGuard Home does not perform cross-list deduplication in memory. When it loads its own native subscription filters (e.g., 'AdGuard DNS filter') alongside our massive 'compiled_rules.txt', overlapping rules are loaded twice, wasting significant RAM in the Colima VM. The UI would show ~500k rules, representing the sum of all lists rather than unique domains.

1911

1912 ****Fix:**** Updated 'scripts/sync-rules.py' to intelligently cross-reference AdGuard's configuration and deduplicate our list against it.

1913 1. The script now reads 'adguard/conf/AdGuardHome.yaml' to dynamically fetch the exact URLs of any enabled native AdGuard filters.

1914 2. The parsing engine was upgraded to normalize basic AdGuard syntax ('||domain') back into raw base domains during the build process.

1915 3. The script subtracts the native AdGuard domains from our custom blocklist *before* compiling, immediately purging ~83,000 completely redundant rules.

1916 4. The normalized base domains then pass through our subdomain optimizer, which squashes an additional ~50% of the remaining rules.

1917

1918 This reduces the final compiled output from ~325k to ~281k rules on the 'heavy' profile, saving ~45k rules from being loaded into memory twice and reducing the overall RAM footprint of the 'adguardhome' container.

1919

1920 ##### 15.12.17 Kernel-Level IP Blocklist (Threat Intelligence Firewall)

1921 **Problem:** DNS sinkholing cannot stop malware or botnets that bypass DNS entirely by hardcoding direct IP addresses (a common technique for C2 communication). These connections pass straight through AdGuard unseen.

1922

1923 **Fix:** Added a second compilation pipeline to 'scripts/sync-rules.py' that runs on every startup alongside the DNS pipeline.

1924 1. The compiler concurrently fetches four curated threat-intelligence feeds: **Feodo Tracker** (abuse.ch botnet C2 IPs), **Spamhaus DROP** (IPs allocated to criminal organizations), **Emerging Threats** (Proofpoint active C2 and scanners), and **CINS** (active brute-force sources).

1925 2. All entries are normalized via Python's 'ipaddress' module. Any IP or CIDR that overlaps with RFC1918 private ranges, loopback, link-local, or the Tailscale CGNAT range ('100.64.0.0/10') is stripped to prevent accidentally locking users out of their own LAN or mesh.

1926 3. The cleaned list (16,721 unique IPs/CIDRs) is written as an 'ipset restore' file using an **atomic swap pattern**: 'create_new populate swap with live destroy_new'. This guarantees the live 'block_malicious' ipset is never empty during a reload.

1927 4. 'scripts/routing-fix.sh' watches the file's mtime every 30 seconds. On change it runs 'ipset restore' and idempotently re-injects 'FORWARD DROP' rules for both 'src' and 'dst' into both iptables backends (nftables + legacy).

1928 5. 'docker-compose.yml' mounts 'adguard/work/userfilters/' into 'routing-fix' as '/userfilters:ro' and adds 'rule-compiler: service_completed_successfully' to its 'depends_on', eliminating the race condition where routing-fix could start before the file existed.

1929

1930 **Memory cost:** The entire 16,721-entry ipset costs only ~1.6 MiB of kernel memory negligible in the 600MB VM budget.

1931

1932 ##### 15.12.18 Incident Post-Mortem: Discarded Docker Exec Errors in logs (July 10, 2026)

1933 **Symptom:** When the 'warp' container is stopped, dead, or failing to connect to its endpoints, the 'toggle.sh' and 'recover.sh' connectivity health checks fail, but no details are printed to 'output.log'. The logs only show generic 'FAIL' outputs, making it hard to see if the failure was a network timeout, Docker daemon error, or a missing binary.

1934

1935 **Root Cause:** Overly broad error silencing. The 'docker compose exec' commands in the health checks (Check 3 in 'toggle.sh' and the traffic check in 'recover.sh') both redirected stderr to '/dev/null' ('>/dev/null' and '2>/dev/null' respectively), completely throwing away any diagnostic error messages produced by Docker or 'wget'.

1936

1937 **Fix (July 10, 2026):** Redirected stderr of the 'docker compose exec' check commands to 'output.log' ('2>> output.log'). This ensures that any command execution failures or connection timeout details are logged.

1938

1939 ##### 15.12.19 Incident Post-Mortem: Toggle Decided STOP/START From a Live Probe, Aborting When Docker Was Down (July 13, 2026)

1940 **Symptom:** Running './toggle.sh' (or '--restart') while the Colima VM / Docker daemon was **not** running caused the script to hang for ~15 seconds and then abort during the START phase without ever booting Colima. With 'DEBUG_TRACE' the lifecycle breadcrumb showed 'toggle.sh EXIT: code=1 phase=start', and the trace showed execution jumping straight from the START-branch entry to the 'ERR' trap ('cleanup_handler ERR') with the START body never executing. Historically this silent failure occurred 57 in 'output.log'. Spam-clicking the toggle during the resulting window (each retry rejected by the lock guard, then finally succeeding) is what produced the earlier "pile of stacked toggle processes" and Wi-Fi-bounce confusion.

1941

1942 **Root Cause:** 'toggle.sh' decided whether to STOP or START by calling 'is_gateway_active()' ('scripts/common.sh'), which **live-probes** the running system its first step is 'run_with_timeout 15 docker compose ps --status running'. When '/var/run/docker.sock' is absent (Colima down), that call blocks for the full 15 s timeout and returns non-zero. Combined with the script's global 'set -e' and 'ERR' trap, a non-zero result evaluated at the 'if is_gateway_active; then else fi' boundary aborted the whole script *before* the 'else' (START) branch body ran precisely the branch whose job is to start Colima. Design-wise the deeper flaw is that a **disable** should not depend on whether the gateway is *currently* healthy; it should act on whether the gateway is *supposed* to be running (persisted intent). The correct signal already existed 'MARKER_FILE=/tmp/nullexit-gateway-active.marker', written at the end of a successful START and cleared on STOP but 'toggle.sh' only wrote/cleared it and never read it for this decision. **Note:** this was a long-standing latent bug, not a regression from the July 2026 'common.sh' platform-function unification or the 'DEBUG_TRACE' work; those changes only made the previously-silent abort *visible* (via the 'EXEC:/'EXIT:' breadcrumbs and xtrace).

1943

1944 **Fix (July 13, 2026):** Introduced 'gateway_should_be_running()' in 'common.sh' and routed all four decision sites through it ('toggle.sh' main decision + '--restart' pre-check; 'scripts/toggle-linux.sh' main decision + '--restart' pre-check). It decides from persisted intent:

1. 'toggle.sh' captures the marker's existence into 'GATEWAY_MARKER_AT_STARTUP' ****before**** the top-of-script "defensive stale-marker clear" wipes it, and the helper honors that captured value (falling back to the live marker file for callers that run before the clear, e.g. the '--restart' pre-check).

2. Marker present STOP path; marker absent START path. This is instant and cannot hang or abort.

3. Only when the marker is absent ****and**** the Docker socket is actually reachable ('/var/run/docker.sock' or '~/colima/default/docker.sock' on macOS) does it consult 'is_gateway_active' as a non-fatal reconciliation fallback (covers a marker lost to a '/tmp' wipe while containers are genuinely up). A dead socket short-circuits to "stopped" with no probe, so 'set -e' can never be tripped.

4. The STOP path's 'docker compose down' is now '|| true'-guarded so a disable always completes even with Docker down; the START path's 'docker compose up' is deliberately left unguarded (a failed start *should* trigger cleanup). On Linux the marker isn't maintained, so the helper degrades to the 'is_gateway_active' fallback fast there because native Docker has no Colima VM to hang.

****How we got here (methodology note).**** This one is worth recording because the **process** mattered as much as the patch. The failure had been happening silently for a long time the log simply ended after a teardown with no error, so there was nothing to chase. We first attacked the ***observability gap*** rather than guessing at the cause: we wrapped the heavyweight lifecycle commands ('colima start', 'docker compose up/down') so their output and exit codes always land in 'output.log', added an 'EXIT: code= phase=' breadcrumb that fires even when the process is killed, and put a full opt-in xtrace behind 'DEBUG_TRACE'. Only **then** did we reproduce the failure and this time the trace said, unambiguously, 'EXIT: code=1 phase=start' with the START body never entered. The instrumentation didn't fix anything, but it converted an invisible abort into a precise, reproducible fact. (Lesson, consistent with 15.12.18 and the LUKS post-mortem referenced in 15.11: ***a mechanism that fails silently is worse than one that fails loudly*** so we make it fail loudly first.)

The root-cause leap, though, came from a ***design-smell intuition, not the stack trace***: the observation that it is nonsensical for a **disable** to first check whether the gateway is **currently online** a teardown should act on whether the gateway is **supposed** to be running, not probe a subsystem to find out. That instinct turned out to name the bug exactly. The smell was a ***conflated responsibility with an inverted dependency***: to decide whether to **start** the gateway, the code first had to **talk to** the gateway (via Docker) the very thing START is responsible for bringing up so the check was guaranteed to be unavailable in precisely the cold-start case that mattered. A second tell was ***false symmetry***: enable (builds state) and disable (tears it down) are not symmetric operations, yet both ran the identical expensive probe; whenever two operations that should differ share suspiciously identical machinery, the code is usually lying about the structure of the problem, and reality eventually calls the bluff. The fix invented nothing the honest signal ('MARKER_FILE', literally "the gateway is supposed to be up") already existed and was maintained in all the right places; it was simply never consulted for the decision it was made for. We just pointed the decision at the answer the codebase already knew. General principle worth keeping: ***design smells here are load-bearing they tend to *be* latent bugs, not merely cosmetic***, and are worth chasing on intuition even before a trace confirms them.

****Acceptance verified (static + live):**** with marker present STOP (instant, even with Docker down); with marker absent + Docker down START (0 s, no 'set -e' abort); a real run now proceeds past init to the gateway-status decision. 'bash -n' clean on all three scripts; no new 'shellcheck' findings; re-signed 'signatures' ('toggle.sh'/'common.sh'/'toggle-linux.sh' are in the signed set) and 'crypto.sh --verify' exits 0.

****Follow-up (same day)** a 'set -e' regression the fix introduced, and why only a live run caught it.****** The first cut of the marker capture wrote the intent as a full 'if [-f "\$MARKER_FILE"]; then GATEWAY_MARKER_AT_STARTUP="yes"; else GATEWAY_MARKER_AT_STARTUP="no"; fi'. On macOS '/bin/bash' 3.2, with 'set -e' active and the 'DEBUG_TRACE' xtrace redirect ('exec 2> >(tee)') in effect, that construct ***aborted* the script at *init*** the moment the marker was absent the false '[-f]' status leaked through the compound and 'set -e' killed the run ('EXIT: code=1 phase=init'), before the gateway logic ran at all. Notably this did ***not*** reproduce in isolated 'bash -c' snippets of the same block it only surfaced in the script's real runtime context so we had to bisect against the actual 'toggle.sh' (swap the block for a trivial assignment the script sailed past init) to prove causation. Fix: write it as a default assignment plus an 'if' ***without* an 'else' **** ('GATEWAY_MARKER_AT_STARTUP="no" then 'if [-f "\$MARKER_FILE"]; then GATEWAY_MARKER_AT_STARTUP="yes"; fi'), which yields status 0 when the condition is false. Lesson reinforced: 'set -e' on legacy bash is genuinely treacherous around conditionals, and the observability work paid for itself again the breadcrumb pinned the abort to init instantly, and a live run (not static checks) was the only thing that exposed a heisenbug invisible to isolated reproduction.

15.12.20 Honey-Port Tripwire Accumulation (July 14, 2026)

****Symptom:**** After many toggle-ons, 'ps' showed a pile of idle 'bash ./toggle.sh' subshells, each blocked on an 'nc -l localhost <port>' listener one leaked per toggle-on, never reaped.

****Root Cause:**** The Honey-Port Tripwire arms a fresh random port each START via a disowned '(nc -l) &' subshell. 'nc -l' blocks until a connection that (by design) almost never comes, so each toggle-on leaks one idle listener; because the port is random, successive listeners never collide and simply stack up over time.

****Fix:**** 'toggle.sh' now reaps the previous tripwire before arming a new one 'pkill -f "nc -l localhost" (macOS) / 'pkill -f "nc -l .*.127.0.0.1" + "nc -l localhost" (Linux) immediately before 'HONEY_PORT=\$(get_free_port)' at both arming sites. No 'sudo' needed (the listeners are user-owned). Net effect: at most one tripwire alive at a time instead of one per toggle. (The same 'pkill' pattern cleared the backlog that this session's ~15 test toggles produced.)

15.12.21 Post-Wake Self-Feedback Loop: The Watcher Re-Fires On The Route Change It Caused (July 14, 2026)

1966 ****Symptom:**** After a container restart (or wake), 'output.log' showed 'recover.sh --post-wake' launching
 ~6 times in a row, one launch roughly every ~40s, each exiting 'exit=0' recovery *worked* every
 time yet kept re-firing "until it worked." Every launch was triggered by 'NET: changedKey State:/
 Network/Global/IPv4'.

1967

1968 ****Root Cause** a self-feedback loop, NOT PF.** 'scripts/watcher.sh' Listener 2 watches 'State:/Network/
 Global/IPv4' via 'scutil n.watch' and launches 'recover.sh --post-wake' on any change. But post-wake
 re-asserts the exit node ('recover.sh' ~L367, 'tailscale set --exit-node'), and even though 'set'
 is deliberately used instead of '--reset' to minimise churn (recover.sh ~L360-366), **re-asserting
 the exit node re-writes the host default route and the default route IS 'State:/Network/Global/IPv4'
 '.* So each run mutates the exact key the watcher triggers on. The self-triggered change lands
 ~15-45s after the run finishes; 'run_recover()'s old 10s debounce (reset to the run's *finish* time
 at the tail of the function) had long expired by then, so the watcher re-fired. The loop self-
 terminates only once 'tailscale set' becomes a no-op (route already correct) no route change no re
 -trigger.

1969

1970 ****Why the debounce couldn't stop it:**** the loop period (~40s = one full post-wake run + settle) far
 exceeds the 10s debounce, which therefore only ever caught macOS's rapid *duplicate* scutil
 notifications for a single change (the '+0.2s' doublets), not the self-trigger.

1971

1972 ****Not PF:**** post-wake never calls 'pfctl' (the kill-switch is armed once, in 'scripts/common.sh', not per
 cycle); no PF-enable line falls inside the loop window. The 'No ALTQ support in kernel' warning is
 normal macOS pfctl noise (Apple ships pfctl without ALTQ compiled in), and 'Could not capture PF
 reference token' is a separate benign ref-count bookkeeping wart both unrelated to the loop.

1973

1974 ****Fix:**** 'scripts/watcher.sh' 'run_recover()' now arms a **post-recovery settle cooldown** ('/tmp/
 nullexit-watcher.settle-until', 'NULLEXIT_POSTWAKE_SETTLE_SECONDS' default 45s) *after* each post-
 wake run, and skips launching while 'now < settle-until'. Because it is armed only after a recovery
 runs, a genuine first wake/roam is still handled immediately; only the Global/IPv4 change the run
 caused itself is absorbed. A real new roam during the window is delayed at most '
 POSTWAKE_SETTLE_SECONDS', an acceptable trade for killing the loop. Watcher-side only 'recover.sh'
 is unchanged. Severity was low (self-correcting, healthy every iteration, no leak/outage; cost was
 wasted ~28s runs + log noise).

1975

1976 **### 15.12.22 FaceTime (and Real-Time P2P Media) Over the Double-Tunnel: Why It Breaks + the Split-Route
 Plan (July 15, 2026)**

1977 ****Symptom:**** A FaceTime call placed while nullexit is up fails to connect.

1978

1979 ****Diagnosis** nullexit is NOT blocking it; every layer passes the traffic (all verified live):**

1980 - ****DNS:**** no Apple/FaceTime domain is sinkholed 'dig' via the gateway ('100.100.21.8') returns real
 Apple A/CNAME records, none '0.0.0.0'/'127.0.0.1'. The '<no answer>' domains ('gateway.push.apple.
 com', 'fmn.apple.com', 'identitieservices.apple.com', 'fdr.apple.com') return empty on '1.1.1.1' too
 CNAME/non-A, not our block.

1981 - ****routing-fix 'FORWARD':**** the exit-path rule '-i tailscale0 -o tun0 -j ACCEPT' carries ~203k pkts / 73
 M bytes and accepts *everything*.

1982 - ****Country / threat blocklists**** ('block_il', 'block_kp', 'block_malicious'): **0** packets dropped.

1983 - ****WARP (gluetun) egress:**** '-A OUTPUT -o tun0 -j ACCEPT' + '-A POSTROUTING -o tun0 -j MASQUERADE' all
 outbound UDP (any port) leaves and is SNAT'd through WARP.

1984

1985 ****Root cause** double symmetric NAT ending at Cloudflare WARP.** Path: 'host Tailscale (MASQUERADE) WARP
 (MASQUERADE) Cloudflare shared egress'. WARP is a CGNAT-style shared egress with **no inbound UDP
 mappings**, and the double-MASQUERADE presents as **symmetric NAT** the exact condition that
 defeats FaceTime's STUN/ICE UDP hole-punching. No P2P path can form; the Apple-relay fallback is
 unreliable through the stacked tunnel, and the 12801200 MTU squeeze degrades media. This is inherent
 to routing real-time P2P media through WARP, **not** a firewall rule. (Zoom/WhatsApp *calls* hit
 the same wall.)

1986

1987 ****Latent finding (separate bug) Closed:**** the 'FORWARD' chain's '-i tailscale0 -p udp --dport 443' (QUIC
) and '--dport 853' (DoT) 'REJECT' rules were **shadowed** by the earlier '-i tailscale0 -o tun0 -j
 ACCEPT' (first-match wins) verified **0 pkts** on each at the time. Fixed by moving those port-
 REJECT' rules in 'post-rules.txt' **before** the 'tailscale0 tun0' / 'tun0 tailscale0' ACCEPT
 rules. Re-verified live post-fix: the REJECTs now sit ahead of the ACCEPTs in 'iptables -L FORWARD
 --line-numbers' and accumulate hits.

1988

1989 ****Fix plan (chosen: Option 2 narrow, always-on, opt-in split-route). Implemented and shipped enabled by
 default:****

1990 - ****Mechanism:**** add more-specific host routes for FaceTime's Apple '/16's via the **physical** gateway
 ('172.17.0.1'/'en0'). Longest-prefix match beats Tailscale's exit-node '0.0.0.0/1' + '128.0.0.0/1'
 routes, so only those '/16's leave **direct** while everything else stays double-tunneled. Mirrors '
 add_warp_bypass_routes'.

1991 - ****Wiring:**** 'add_apple_split_routes()' / 'remove_apple_split_routes()' in 'common.sh'; called in toggle
 START (after exit-node assert) + 'recover.sh --post-wake' (survives roam) + removed on STOP;
 guarded by 'FACETIME_SPLIT_ROUTE' in '.env' '.env.example' ships it 'false' (opt-in, privacy-first
 default for new installs), enabled here after the tradeoff below was accepted and the CIDR verified
 against a real call.

1992 - ****Privacy tradeoff:**** whatever '/16's go direct see the **real ISP IP** (not WARP). Apple already ties
 traffic to the Apple ID, so Apple-direct is the accepted scope; user chose **narrow** (FaceTime-
 infra '/16's only) over all-'17.0.0.0/8' to minimise the surface.

1993 - ****CIDR list finalised from an empirical capture:**** host-side 'sudo tcpdump -n -i <exit-utun, e.g.
 utun6> net 17.0.0.0/8' during a real FaceTime call. The capture identified **'17.249.0.0/16'** as
 the load-bearing media/relay range this is the value shipped in '.env.example's '

FACETIME_DIRECT_SUBNETS', verified working end-to-end on a real call. (Earlier DNS-only guesses '17.57/16' APNs push/ringing, '17.248/16' iCloud/FaceTime gateway, '17.253/16' aapling push CDN, '17.157/16' GSA auth were signaling-only ranges and did not carry media; the capture superseded them.) See P2P note 16.4.

15.12.23 setup.sh Never Generated NULLEXIT_SEED Every Fresh Install Hard-Failed Its First toggle.sh (July 18, 2026)

****Symptom:**** README "Cryptographic Integrity Verification" and this doc's own '.env' variable table both claimed 'NULLEXIT_SEED' is "generated by 'setup.sh'." It wasn't. 'scripts/setup-common.sh's '.env'-writing step never referenced 'NULLEXIT_SEED' at all, and '.env.example' only offered a placeholder telling the user to run 'openssl rand -hex 32' by hand. '.signatures' is committed to git (tracking the maintainer's own local seed), so a fresh clone's 'toggle.sh' would call 'scripts/crypto.sh --verify', find no 'NULLEXIT_SEED' in the new '.env' at all, and hard-fail with a cryptic seed error on the very first run with nothing in the setup flow or its closing instructions telling the user a 'crypto.sh --sign' step was needed, or why.

****Compounding bug found in the same code path:**** 'setup-common.sh's '.env' writer also defaulted 'KILL_SWITCH' to "false" for brand-new installs ('EXISTING_KILL="false"', only overridden if an '.env' already existed with a different value) the exact inverse of the "leave true, this is the core leak guarantee" invariant enforced everywhere else ('.env.example', 'scripts/profiles.sh'). Every fresh install would have shipped with the kill-switch off by default, not just failed on first boot.

****Fix:**** 'setup-common.sh' now generates 'NULLEXIT_SEED' via 'openssl rand -hex 32' when '.env' doesn't already have one (preserving it on re-runs, same pattern as every other preserved flag), writes it into the new '.env', and calls 'scripts/crypto.sh --sign' immediately after writing '.env' so a fresh install's first 'toggle.sh' finds a valid seed and matching signatures instead of failing cold. 'EXISTING_KILL' now defaults to "true".

Part IV Observations, Changelog & TODO

16. Unverified Observations

16.1 AdGuard Returns Two Different Sinkhole IPs (Unverified)

> ****Status:** Observed but not formally verified. Needs a deeper audit of the rule-compiler sources to confirm.**

****Observation (July 10, 2026):**** When querying AdGuard Home for blocked ad domains, some domains resolve to '0.0.0.0' and others resolve to '127.0.0.1'. Both are blocked, both cause connection failure, but the different responses were unexpected.

doubleclick.net	0.0.0.0	(blocked)
ads.google.com	127.0.0.1	(blocked)
pagead2.googlesyndication.com	0.0.0.0	(blocked)
scorecardresearch.com	127.0.0.1	(blocked)

****Likely Explanation (Unverified):**** There are three historical schools of thought on the "correct" DNS sinkhole response, and AdGuard faithfully preserves whichever the source blocklist used:

Sinkhole IP	Convention	Origin
'0.0.0.0'	AdBlock-style lists ('\\ domain~')	Modern DNS-level blocking; AdGuard's native format. Fails fast OS doesn't attempt a TCP connection. Works for both IPv4 and IPv6 without a separate rule.
'127.0.0.1'	Hosts-file-style lists ('127.0.0.1 domain')	90s-era '/etc/hosts' tradition. Steven Black's hosts list (500k+ domains) still uses this.
'NXDOMAIN'	Purist / Pi-hole default	Returns "domain doesn't exist." Semantically most accurate but requires browsers to handle NXDOMAIN gracefully.

The dual-response behavior in this project is because 'rule-compiler' pulls from both AdBlock-format sources (Hagezi, uBlock origin lists) and hosts-format sources (Spamhaus DROP, Steven Black, Feodo Tracker), and AdGuard returns whatever sinkhole IP the matching rule specifies.

****To Verify:**** Run 'docker compose exec tailscale cat /etc/hosts' to confirm no hosts-file rules are being injected at the container level, and check which specific AdGuard rule matched each domain via the AdGuard query log at 'http://100.100.21.8:3000/#logs'.

16.2 AGY CLI Appears to Generate Native macOS Notification Pop-ups (Unverified)

> ****Status:** Observed but source unverified. macOS security logs ruled out system-level origin.**

****Observation (July 10, 2026):**** While running a DNS blocklist verification via the Antigravity CLI ('agy'), a command was proposed that included 'malware.wicar.org' (a known-safe DNS blocklist test domain) in a 'dig' loop. Shortly after, a macOS-style notification popup appeared describing a "malicious" or "blocked" event. The AGY CLI terminal session then closed.

****Investigation:****

2040 - **macOS XProtect / Gatekeeper logs:** Empty. Zero events in the 30-minute window around the incident.
 2041 - **macOS Endpoint Security / System Policy logs:** Empty.
 2042 - **Broad 'log show' search for "malicious", "malware", "blocked":** Empty.
 2043 - **Conclusion:** macOS itself did not generate the notification. No system security subsystem fired.
 2044
 2045 **Likely Explanation (Unverified):** The notification appears to have originated from **AGY CLI's** own internal safety system. AGY CLI is a native macOS application (not just a terminal process) and has access to the macOS Notification Center API ('UNUserNotificationCenter'). When its guardrails detected a domain containing the word "malware" ('malware.wicar.org') in a proposed shell command, it likely:
 2046 1. Blocked the command from executing
 2047 2. Posted a native macOS-style notification via the Notification Center
 2048
 2049 This is surprising behavior for a CLI tool most terminal programs don't send GUI notifications but AGY CLI appears to bridge both worlds: it runs in a terminal but is packaged as a native app with full system API access. The notification would be indistinguishable from a system security alert at a glance.
 2050
 2051 **Notes:**
 2052 - 'malware.wicar.org' is the DNS/AV equivalent of the EICAR test file a deliberately benign domain that triggers blacklist/AV systems to verify they work. It was included in the test set to confirm the AdGuard blacklist was catching it. AGY's safety system is not aware of this distinction.
 2053 - For future DNS blacklist testing, use only ad/tracking domains (e.g., 'doubleclick.net', 'adnxs.com') to avoid triggering AGY's guardrails.
 2054
 2055 **## 16.3 'system_profiler SPAirPortDataType' Longevity (Wi-Fi Detection)**
 2056 The Wi-Fi security detection used by 'watcher.sh's 'detect_lan_p2p_mode()' currently relies on 'system_profiler SPAirPortDataType' (after the 'airport' binary removal full history in 15.11.6). This works as of macOS Sequoia (15.x) without sudo and without redaction, but given Apple's ongoing Wi-Fi privacy clampdown, it may be locked down in a future release. If 'security' comes back empty on a future macOS version while on Wi-Fi, detection silently falls back to 'allow=false' (a safe default). The long-term fix would be a small Swift CoreWLAN helper compiled and bundled in the repo.
 2057
 2058 **## 16.4 Tailscale LAN P2P Succeeds Cross-AP But Fails Same-AP (Client Isolation + No NAT Hairpin)**
 2059
 2060 **Observation (July 15, 2026):** On the residence Wi-Fi, a Galaxy S24 in the lobby reception 6 floors away, associated to a **different** AP establishes a **direct**, low-ping **Tailscale P2P** session to an iPhone 11. The **same** pair on the **same** AP / same floor / close range does **not** go direct. Counterintuitively, **farther** apart = P2P works. The effect shows for the iPhone 11 (a plain client) but **not** for S24macOS or S24the gateway container.
 2061
 2062 **Mechanism:**
 2063 - **Same AP:** the AP enforces **client isolation** (L2 device/device block on one radio). Both devices then share one public IP; STUN returns identical external mappings, so reaching each other would need the residence router to **NAT-hairpin** which most campus/consumer gear does not do direct P2P fails and falls back to a DERP relay.
 2064 - **Different AP (floors away):** large residence networks **segment** per AP/area into distinct VLANs/subnets, so the two devices sit on different L3 networks routed through the campus core, where client isolation does not apply and the STUN path works **without** hairpin clean direct P2P.
 2065 - **Why macOS / the container are immune:** **nullexit** pins a **direct WireGuard endpoint** (Tailscale port '41642', 'direct 172.x'), so those peers already hold a negotiated direct path and do not depend on same-AP L2 discovery. A plain iPhone 11 has no pinned endpoint, so its path is entirely at the mercy of the Wi-Fi topology which is exactly why **it** exposes the isolation effect.
 2066
 2067 **Relevance:** refines 'watcher.sh's 'detect_lan_p2p_mode()' (which probes same-subnet AP isolation via a broadcast ping). The probe correctly flags same-AP isolation, but this observation adds that isolation is **per-AP** and cross-VLAN P2P can succeed even when same-AP fails so a single "LAN P2P allowed?" boolean under-describes a multi-AP network. No code change; documents expected behavior. Related: the FaceTime NAT analysis in 15.12.22 (same STUN/hairpin physics, different symptom).
 2068
 2069 **## 16.5 Working Theory iPhone Cross-Device Visibility Reaches Other Phones, Not macOS (Instagram Auto-Open, July 15, 2026)**
 2070
 2071 **Observation:** Opening Instagram on a Galaxy S24 (specifically via **Chrome** on the S24, not Samsung Internet) auto-opens a tab on a nearby MacBook, but **only** when the S24 is physically close to the Mac it does not happen when the phone is away. **Chrome** on the S24 has zero granted Android permissions (no Bluetooth, no location, no nearby-devices) this rules out Chrome itself performing any BLE/Nearby proximity scan, since Android enforces those permissions at the app level regardless of what the browser ships. This pushes the likely actor down to **Google Play Services**, which holds system-level Bluetooth/nearby-device access independent of per-app grants and could be doing the proximity detection on Chrome's behalf (e.g. via Android's "Nearby" APIs) then handing the result to Chrome only for the actual tab-open. A purely local, short-range channel is still required to explain the proximity gating regardless of which component performs it; this also rules out any cloud-mediated mechanism (Chrome cloud tab sync, FCM push, Tailscale/WAN P2P per 16.4) since those are distance-independent. Live probes during the session ('system_profiler SPBluetoothDataType' for a paired device, 'dns-sd -B _services._dns-sd._udp' for 8s) found **no** paired Bluetooth device and no Android/Google/Instagram mDNS service **only** the Mac's own Bonjour services (AirPlay, RFB, '_companion-link', 'Spotify Connect'). Inconclusive: a BLE-advertisement-only channel (no formal pairing) would not show up in either probe.

2073 ****User's working theory:**** the iPhone (also on hand, on the same network/Apple ID as the Mac) makes * other devices connected to it* even an Android phone visible to nearby devices, but this visibility channel reaches **other phones more readily than it reaches macOS**. I.e., the iPhone may be acting as a discovery bridge/relay between the S24 and the Mac rather than the S24 and Mac talking directly. Combined with the Play Services angle above, a refined version is: ****Play Services on the S24 detects proximity to an Apple mesh (iPhone and/or Mac both signed into iCloud/AWDL-visible) via some cross-ecosystem discovery Google has built**, rather than the S24 and Mac negotiating directly.**

2074

2075 ****Cross-ecosystem refinement:**** this isn't Google-only. Apple runs an equivalent OS-level background BLE scanning layer for its own Continuity stack (Handoff, Universal Clipboard, AirDrop discovery, Find My the latter using *every nearby Apple device*, not just the owner's, as a passive BLE relay), implemented as system daemons ('bluetoothd'/'sharingd'/'rapportd', already implicated in the unrelated AWDL wedge bug at 15.10) rather than anything gated by app-level permission. The key mechanism that makes the two ecosystems mutually visible: ****BLE advertisement packets are broadcast in the clear over open air**** Apple's Continuity BLE frames are semi-proprietary but unencrypted at the advertisement layer, so any BLE radio (including Android's, via Play Services once Android's own "Wi-Fi & Bluetooth scanning" system-level toggle is on) can passively observe "an Apple device is here" and react, without pairing or explicit permission from Apple's side. So both ecosystems maintain a BLE mesh-discovery layer that sits below their own per-app permission model, and each side's layer is passively readable by the other's radio which is a plausible bridge for a Play-Services-mediated Chrome trigger reacting to iPhone/Mac Continuity broadcasts.

2076

2077 ****Status:**** Unverified. Not yet reproduced under an active capture the July 15 session ran the mDNS browse *before* confirming the phone was actively broadcasting, so it's not a clean negative result. To confirm: capture 'tcpdump -i en0 udp port 5353' (mDNS) ****and a raw BLE advertisement capture**** (e.g. 'nRF Connect' or similar BLE scanner app on a third device, or 'sudo btmon'-equivalent on Linux macOS does not expose raw BLE advertisement capture without a hardware sniffer) at the exact moment Instagram is opened on the S24 while close, with the iPhone present vs. absent, to isolate whether the iPhone's Continuity broadcasts are the trigger or the S24Mac channel exists independently of it. On the Android side, checking 'Settings Location Location Services Wi-Fi & Bluetooth scanning' and whether ****Google Play Services**** (not Chrome) holds Bluetooth/nearby-device/location permissions would confirm or kill the Play-Services-as-actor theory. Related: 16.4 (same physical-proximity family of P2P/discovery quirks, different mechanism and inverted distance relationship).

2078

2079 **## 17. Changelog / Recent Updates**

2080

2081 Reflects the current branch's state. Each entry is a one-line summary + commit hash; read the commit message (and the linked Incident Log entry) for full detail. Full post-mortems live in 15.

2082

2083 - ****July 10, 2026 'fix(race): suppress WARP Watcher during post-wake force-recreate**** Fixed a race where switching Wi-Fi caused the host IP to leak as the raw ISP IP ('132.x') on every roam; 'recover.sh --post-wake''s warp force-recreate tripped the WARP Watcher into a nuclear shutdown. Fixed via '/tmp/nullexit-warp-inhibit.marker'. See 15.2.7.

2084 - ****July 10, 2026 startup/roaming stabilization suite**** Pre-flight Tailscale race (15.5.3), STUN-block cold-boot false negative (15.5.4), restart DNS-build race via 'wait_for_dhcp_settle' (15.5.5), nuclear/post-wake concurrency lock (15.2.8), infinite auto-recovery loop marker leak (15.2.9), PF kill-switch bricking SOCKS5 fallback (15.7.3), '--reset' host logouts (15.7.4), missing-flags abort (15.7.5), discarded docker-exec errors (15.12.18), and the SNAT endpoint poisoning fix (15.7.1).

2085 - ****July 6, 2026 'fix: resolve edge-case vulnerabilities and system state leaks**** Security-review hardening suite:

2086 1. ****DNS Watcher Interface Independence:**** 'start_dns_watcher' now detects the active interface via 'get_active_service()' instead of hardcoding a 'grep' for "Wi-Fi".

2087 2. ****Robust Lock Verification:**** 'toggle.sh' lock-file logic uses 'ps -p' verification to prevent permanent lockouts from a recycled PID.

2088 3. ****Fail-Closed Crypto:**** 'crypto.sh' integrity check switched from '[-x]' to '[-f]' so it runs even if git strips the '+x' bit.

2089 4. ****Strict 'iptables' Pinning:**** 'post-rules.txt' FORWARD-chain rules explicitly pin 'tailscale0'/'tun0' both directions, preventing escape via 'eth0' during WARP drops.

2090 5. ****Docker Local Network Isolation:**** 'docker-compose.yml' binds exposed WARP ports ('1080', '5354') to '127.0.0.1' to prevent LAN access.

2091 6. ****Routing Verification & Sudoers:**** Added a 'netstat -nr' default-route override check and ensured '/usr/bin/python3' is permitted in the sudoers template.

2092 - ****July 4, 2026 'cc789c5**** Concurrency mutex on 'toggle.sh' ('/tmp/nullexit-toggle.lock'); defensive 'TS_AUTHKEY' init under 'set -u'; Go 'go vet'/'go build' steps in CI; AdGuard log capping; macOS BSD 'grep' fixes in 'setup-common.sh'; 'recover-linux.sh' quoting/PID fixes; removed redundant bypass-routes block in 'recover.sh'.

2093 - ****July 4, 2026 'fix: post-wake routing loop & destructive recovery**** (1) 'add_warp_bypass_routes' now accepts explicit interface args so post-wake stops routing WARP traffic back into 'utun*' (was an infinite loop killing internet on every wake). (2) Replaced 'sudo route -n flush' with targeted deletions of '0.0.0.0/1'/'128.0.0.0/1' + WARP bypass routes, so recovery no longer destroys loopback/multicast routes and wedges 'configd'/'mDNSResponder' until reboot.

2094 - ****July 4, 2026 'feat: secure execution**** Restricted the blanket '/usr/bin/python3' NOPASSWD sudo to exactly '/usr/bin/python3 -I .../scripts/dns-proxy.py'; locked 'dns-proxy.py' with 'chown root:wheel' + 'chmod 755'; added native auth prompts to 'Toggle Gateway.app'/'Recover Gateway.app'; 'toggle.sh' now captures 'warp' logs on failure before teardown.

2095 - ****July 4, 2026 Go ports + refactor**** 'logger' and 'rule-compiler' ported to Go multi-stage Docker builds (15.8.6); ~200 lines of duplicated bash extracted to 'scripts/common.sh' (15.9.4); 'crypto.sh' HMAC-SHA256 script integrity added.

```

2096 - **July 3, 2026 'fix: resolve numerous system regressions'** Fixed truncated '/etc/sudoers.d/nullexit'
    line; temporary 'tmp'+os.replace()' atomic writes in 'sync-rules.py' (later reverted, see below);
    'recover.sh' 'ts_args' quoting; ANSI '%b'/'-e' output in 'diagnose-host-leak.sh'/'fix-docker-bridge
    -collision.sh'; 'toggle-linux.sh' using 'resolvectl dns' instead of macOS 'networksetup'; AdGuard
    REST refresh POST targeting port '3000'; restored 'START_GATEWAY=true'; purged stale port '80'
    mapping; verified 'output.log' in '.gitignore'.
2097 - **July 3, 2026 'unlock-files.sh' + revert atomic writes** Reverted all 5 'tmp'+os.replace()' sites
    in 'sync-rules.py' to direct writes; created 'scripts/unlock-files.sh' to permanently fix stale
    '0444'/'000' file permissions via atomic rename (no 'chmod'). Covers '.env', 'compiled_rules.txt', '
    ip_blocklist.ipset', 'cache/*.txt', 'cache/ip/*.txt', 'data/filters/*.txt'. See 15.9.6 + 3.
2098 - **July 3, 2026 Overnight leak + geo-blocking** Introduced 'scripts/host-leak-probe.sh' (sub-second
    host-egress detector, 15.6.1) and the statistical-threshold WARP auto-shutdown watcher (15.6.2);
    fixed the overnight silent IP leak (15.6.3).
2099 - **July 2, 2026 '8c5fae1' + '181e1e1'** Two infinite routing loops fixed: host-VM tunnel loop (WARP
    bypass routes via physical gateway IP + 'setup_exit_node_routing' forcing default to 'utun*') and
    Docker Compose subnet takeover ('10.200.1.0/24' lock). '--accept-routes=true' added explicitly to
    all 'tailscale up --reset' calls. See 15.2.2.
2100 - **July 2, 2026 auto-recovery daemon** 'scripts/watcher.sh' + launchd LaunchAgent documented. See
    15.5.1.
2101 - **July 1, 2026 'c5b92c0' (refactor: personal-leak cleanup)** Eliminated every residual personal-
    username leak across source + config; launchd plist uses 'WorkingDirectory' + relative program arg
    with the '__NULLEXIT_HOME__' sentinel (install recipe gains a single 'sed -i ' 's|__NULLEXIT_HOME__
    |$(pwd)|' step); 8 devref prose sites genericized to '~/colima/...', '/Library/LaunchAgents/...',
    '$USER', '<USER>'.
2102 - **July 1, 2026 'a5e2a5a' (refactor: script-relative paths)** Replaced 6 hardcoded '/Users/<user>/...'
    nullexit/... literals with script-relative resolution. 'recover.sh'/'watcher.sh' inject 'SCRIPT_DIR
    ='${cd "$(dirname "${BASH_SOURCE[0]}")" && pwd}"; 'RECOVER="${SCRIPT_DIR}/../recover.sh"' resolves
    without launch-path dependency.
2103 - **June 28, 2026 '8f8e4e' (M1: scripts/ move)** 8 internal scripts consolidated into 'scripts/'; the
    4 user-facing/orchestrator/config files ('toggle.sh', 'recover.sh', 'setup.sh', 'docker-compose.yml
    ') stayed at repo root.
2104 - **June 28, 2026 '0875ef3' (ci: glob)** CI 'py_compile' switched to 'python3 -m py_compile scripts/*.
    py' after M1 left root with zero '.py' files.
2105 - **June 28, 2026 '25b37a1' (docs: scripts/ prefix)** 'devref.md' + 'README.md' updated to the 'scripts
    /' prefix for all 8 moved files ('31 patches').
2106
2107 ### End-to-end verification (commit 'c5b92c0')
2108 Manual START STOP cycle ran clean on macOS after path relativization + personal-leak cleanup:
2109 - 'bash toggle.sh' START: exit 0, ~135s (cold Colima boot); marker present ('2026-07-02T03:41:41Z'), all
    5 containers healthy, host DNS hijacked to '100.100.21.8' (no fallback), exit node selected,
    Cloudflare trace through WARP succeeds.
2110 - 'bash toggle.sh' STOP: exit 0, ~12s; marker cleared, all nullexit containers gone, host DNS restored to
    '1.1.1.1', Tailscale disconnected.
2111
2112 ## 18. TODO & Future Work
2113
2114 ### 18.1 TODO
2115 * **FaceTime / real-time P2P split-route (Option 2, narrow)** ~Implement 'add_apple_split_routes()' / '
    remove_apple_split_routes()' in 'common.sh' (route FaceTime's Apple '/16's direct via the physical
    gateway, opt-in '.env' flag, default OFF), wired into toggle START/STOP + 'recover.sh --post-wake'.
    **Blocked on** an empirical 'tcpdump' capture of the relay/STUN '/16's during a real call (user-
    triggered).~~ **Closed** implemented and shipped enabled by default. The 'tcpdump' capture from a
    real FaceTime call confirmed '17.249.0.0/16' as the load-bearing media/relay range; '
    FACETIME_DIRECT_SUBNETS' in '.env.example' ships that value. Full design + diagnosis in 15.12.22.
2116 * **Fix the shadowed exit-path port REJECTs (DoT/QUIC bypass)** ~The 'FORWARD' '-i tailscale0 -p udp --
    dport 853'/'--dport 443' REJECT rules are shadowed by the earlier 'tailscale0 tun0 ACCEPT (0 pkts)
    ', so the intended block-DoT / force-AdGuard and QUIC-block are **not enforced** a mesh client can
    bypass AdGuard via DoT (853). Re-order the REJECTs before the ACCEPT (or scope '-o tun0').~~ **
    Closed** moved the REJECT rules in 'post-rules.txt' before the 'tailscale0 tun0' / 'tun0
    tailscale0' ACCEPT rules. Live 'iptables -L FORWARD -n -v --line-numbers' now shows the REJECTs
    ahead of the ACCEPTs and actually accumulating hits. See 15.12.22.
2117 * **Decide toggle STOP/START from persisted *intent*, not a live probe:** ~toggle.sh' decided STOP-vs-
    START by calling 'is_gateway_active()', which live-probes the running system ('docker compose ps' +
    DNS + SOCKS checks). When Colima/Docker was down the 'docker compose ps' blocked for the full 15 s '
    run_with_timeout' window and, under 'set -e' + the 'ERR' trap, aborted the script *before the START
    body ran* so the path that would boot Colima never executed ('EXIT: code=1 phase=start', seen 57
    historically).~~ **Closed** added 'gateway_should_be_running()' in 'common.sh', which reads
    persisted intent ('MARKER_FILE') captured into 'GATEWAY_MARKER_AT_STARTUP' before the defensive
    stale-marker clear, and only falls back to 'is_gateway_active' when the marker is absent *and* the
    Docker socket is actually reachable (so a dead-socket probe can neither hang nor trip 'set -e'). All
    four decision sites ('toggle.sh' main + '--restart', 'toggle-linux.sh' main + '--restart') now use
    it, and the STOP 'docker compose down' is '|| true'-guarded so a disable always completes even with
    Docker down. See 15.12 Misc Resolved.
2118 * **Verify Wi-Fi Roaming:** Investigate and test whether switching Wi-Fi networks now successfully and
    smoothly recovers the gateway end-to-end without any tailscale flag errors or dead tunnels. (Pending
    user verification).
2119 * **Fix Diagnostic Script Check 5/8:** ~Investigate why 'diagnose-host-leak.sh' check 5/8 ('Host default
    route') sometimes reports "default route goes via physical Wi-Fi Tailscale route assertion failed"
    on macOS even though the 'warp=on' egress checks confirm that traffic is successfully tunneling.~~
    **Closed** fixed by switching check 5 from 'netstat -rn | grep default' to 'route -n get 1.1.1.1'.
    macOS Tailscale uses interface-scoped routes and does not replace the DHCP default route entry. See

```

15.7.1 Known Unknowns.

2120 * **Expose Tor Control Port (9051):** Consider mapping the Tor Control port to localhost and adding 'nyx' (the Tor terminal status monitor) to allow users to inspect their entry, middle, and exit node IPs in real-time. **Closed** mapped the Tor Control Port to a procedurally generated, localhost-bound port 'TOR_CONTROL_PORT', configured 'ControlPort' and disabled cookie authentication in 'docker/tor/entrypoint.sh', and integrated it into the '/sweep' verification protocol.

2121 * **Host-Only Tor Mode (Pluggable Egress):** Implement an 'EGRESS_TYPE=tor_host_only' mode for users in heavily censored (DPI-restricted) environments. This mode would skip booting 'warp' and 'tailscale', boot 'adguardhome' and 'tor' (using Telegram 'obfs4' bridges), set AdGuard's upstream to Tor's 'DNSPort', and use the 'pf' kill-switch to force all host Mac traffic strictly through the Tor network. This provides a 100% free, DPI-resistant evasion path without requiring an external VPS.

2122 * **The Technical Realities (What You Need to Watch Out For):**

2123 * **The UDP Problem:** Tor only transports TCP traffic. It does not support UDP (aside from DNS requests passed directly to its DNSPort). macOS is incredibly chatty over UDP (FaceTime, mDNS, QUIC protocols). Your pf rules defined in 'scripts/pf.conf' must be configured to aggressively and silently DROP all host UDP traffic (except for local DNS queries to AdGuard). If you don't drop it cleanly, macOS services might hang indefinitely while waiting for timeouts.

2124 * **The "Daily Driver" Friction:** Routing a whole host OS through Tor is brutal on performance. Background daemons (iCloud sync, macOS software updates, Spotlight indexing) will blindly attempt to download gigabytes of data over Tor, which could saturate the circuits or time out completely.

2125 * **Bridge Rot:** State firewalls actively hunt and block obfs4 bridges. When a user's bridges burn out, they will lose all internet access due to the pf kill-switch. You will need to ensure the user has a frictionless way to update the 'tor-bridges.txt' file and restart the Tor container without exposing their IP.

2126 * **Exit Node Discrimination:** Because all traffic will exit from known Tor nodes, the user will face an onslaught of Cloudflare CAPTCHAs, and many banking or streaming services will flat-out reject the connections.

2127

2128 **## 18.2 Future Work**

2129 - **Direct P2P on VM Hosts:** Non-Linux hosts run Docker inside a VM, making true UDP hole-punching to arbitrary external peers impossible; those peers fall back to the DERP relay unless bridged mode is enabled on a trusted LAN. The only fully general solution is a native Linux host (Raspberry Pi, Intel NUC) where Docker runs without VM hypervisor translation. *(The one case that works regardless is direct **hostcontainer** P2P the gateway and its own Mac, the same physical machine re-permitted via routing-fix 'host_ips' RETURN rules and confirmed 'direct' even in plain user-mode networking; see 15.3.5 and 15.8.7.)*

2130 - **Native Linux Deployment:** Benchmark on a Raspberry Pi native container footprint is ~75MB without macOS hypervisor overhead.

2131 - **Mesh-Wide Filesystem (SFTP over Tailscale):** Any mesh device passively exposes its filesystem over SFTP. Android via Termux + OpenSSH + wakelock; iOS is client-only due to background process limits.

2132 - **Post-Quantum Cryptography (PQC):** Tailscale uses Curve25519, theoretically vulnerable to "harvest now, decrypt later" quantum attacks. A future iteration could replace the Tailscale container with raw WireGuard + [Rosenpass](https://rosenpass.eu/) to negotiate post-quantum PSKs.

2133 - **Egress Compartmentalization (Pluggable Architecture):** See 12.9 for the full plan to make the entire egress layer swappable via a single '.env' variable.

5 diagrams.md

```

1 # nullexit Flow Diagrams & Architecture
2
3 ---
4
5 ## 1. System Architecture What's Running & Where
6
7 ``mermaid
8 graph TB
9     subgraph INTERNET[" Internet / Cloudflare Edge"]
10         CF["Cloudflare WARP\n(Exit Point)\nwarp=on"]
11     end
12
13     subgraph TAILNET[" Tailscale Mesh (100.x.x.x)"]
14         PHONE[" Phone / Laptop\n(Tailnet Client)"]
15     end
16
17     subgraph MACHOST[" macOS Host"]
18         direction TB
19         PFWALL["macOS pf Firewall\n(Kill-Switch & TCP MSS Clamping)"]
20         TSDAEMON["tailscaled\n(brew service)\nhost Tailscale"]
21         HOSTDNS["Host DNS\nGateway IP\n(hijacked by toggle.sh)"]
22         HOSTSOCKS["Host SOCKS5\n127.0.0.1:1080\n(fallback only)"]
23         HOSTTOR["Host Tor SOCKS5\n127.0.0.1:$TOR SOCKS_PORT\n(randomized)"]
24         HOSTTORCTRL["Host Tor Control\n127.0.0.1:$TOR_CONTROL_PORT\n(randomized)"]
25
26         subgraph MONITORS[" Background Daemons (toggle.sh children)"]
27             WARPWATCHER["WARP Watcher\nevery 30s, via docker exec\noutput.log"]
28             DNSWATCHER["DNS Watcher\n(re-hijacks if drift detected)\noutput.log"]
29         end
30     end

```

```

29     CAFFEINATE["caffeinate -i\n(sleep prevention)"]
30     end
31
32     subgraph COLIMAVM[" Colima VM (Linux, 600MB RAM)"]
33         subgraph NETNS["warp container network namespace (shared by ALL containers)"]
34             GLUETUN["warp / Gluetun\nWireGuard tun0\n(owns the namespace)"]
35             TS["tailscale container\nAdvertises exit node\ntailscale0 interface"]
36             SOCKS["tailscale\nSOCKS5\nport 1080"]
37             ADGUARD["adguardhome\nDNS sinkhole\nport 5335"]
38             ROUTINGFIX["routing-fix sidecar\nGeo-IP Blocking & Go Logger\n(writes blocked.log)"]
39             TOR["tor container\nSOCKS5: port 9050\nControl: port 9051\nTransparent Proxy: port 9040\n
nDNSPort: port 5353"]
40         end
41     end
42
43     subgraph LAUNCHD[" launchd (always-on)"]
44         WATCHER["scripts/watcher.sh\n(sleep/wake + Wi-Fi roam\n+ LAN P2P auto-detection)"]
45     end
46 end
47
48 subgraph RECOVERY[" Recovery"]
49     RECOVER["recover.sh\n(nuclear reset)"]
50     RECOVERPOST["recover.sh --post-wake\n(light refresh)"]
51 end
52
53 %% Traffic flow
54 PHONE -->|"WireGuard / Tailscale mesh"| TS
55 TS -->|"exit-node route tun0"| TS
56 TS -->|"FORWARD tun0 (MASQUERADE)"| GLUETUN
57 GLUETUN -->|"WireGuard UDP (macOS Host Egress)"| PFFIREWALL
58 PFFIREWALL -->|"TCP MSS Clamped / Kill-Switch check"| CF
59 TS -->|"198.18.0.0/15 PREROUTING redirect"| TOR
60 TOR -->|"internal path / exit nodes via tun0"| GLUETUN
61
62 %% DNS
63 HOSTDNS -->|"UDP:53 TCP:5354"| ADGUARD
64 ADGUARD -->|"DNS upstream via tun0"| CF
65 ADGUARD -->|"onion queries localhost:5353"| TOR
66
67 %% Monitoring
68 WARPWATCH -->|"docker exec warp wget cdn-cgi/trace"| GLUETUN
69 DNSWATCHER -->|"networksetup -getdnsservers"| HOSTDNS
70
71 %% Logging
72 ROUTINGFIX -->|"writes"| BLOCKEDLOG["blocked.log\n(Host-side)"]
73 HOSTTOR -->|"port mapping"| TOR
74 HOSTTORCTRL -->|"port mapping"| TOR
75
76 %% Recovery triggers
77 WARPWATCH -->|"6 consecutive warp=off"| RECOVER
78 WATCHER -->|"sleep/wake / roam"| RECOVERPOST
79 RECOVERPOST -->|"if gateway unhealthy"| RECOVER
80 WATCHER -->|"writes .lan_p2p_detected"| ROUTINGFIX
81
82 style MONITORS fill:#1a1a2e,color:#e0e0ff
83 style COLIMAVM fill:#0d2137,color:#e0e0ff
84 style NETNS fill:#0a1628,color:#e0e0ff
85 style RECOVERY fill:#2d0a0a,color:#ffe0e0
86 style INTERNET fill:#0a2d1a,color:#e0ffe0
87 style TAILNET fill:#1a2d0a,color:#e8ffe0
88 style LAUNCHD fill:#2d1a2d,color:#ffe0ff
89
90
91 ---
92
93 ## 2. toggle.sh Full START Flow
94
95 mermaid
96 flowchart TD
97     START(["./toggle.sh"])
98     CHECK{"is_gateway_active?\n(containers running OR\nDNS 1.1.1.1)"}
99
100     START --> CHECK
101     CHECK -->|"YES already ON"| STOPFLOW[" See STOP Flow"]
102     CHECK -->|"NO start it"| S1
103
104     S1["1. Reset DNS 1.1.1.1\n(prevent deadlocks)"]
105     S2["2. tailscale down\n(prevent exit-node routing during boot)"]
106     S3["3. Check Colima VM\n(colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged)"]
107     S4["4. Configure VM swap\n(400MB, prevent OOM)"]

```



```

108 S5["5. Clean corrupted AdGuard config\n(prevent container crash loop)"]
109 S6["6. docker compose up -d --build\n(compile Go threat rules & boot containers)"]
110 S7["7. Poll tailscale status\n(wait for 'offers exit node'\nup to 60s)"]
111 TSREADY{"container\nTailscale ready?"}
112 ABORT([" ABORT\n cleanup_handler"])
113 S8["8. Resolve gateway Tailscale IP\n(.gateway_ip or docker exec)"]
114 S9["9. Verify tailscaled running\n(auto-start if needed)"]
115 S10["10. tailscale up --reset\n--exit-node=\n--accept-routes=true\n--ssh=true --accept-dns=false"]
116
117 PF1{"Pre-flight 1/3\ntailscale ping gateway"}
118 PF2{"Pre-flight 2/3\ndig +tcp AdGuard DNS"}
119 PF3{"Pre-flight 3/3\nWARP internet check"}
120 ALLPASS{"All 3 pass?"}
121
122 EXITPATH[" EXIT NODE PATH\n tailscale up --exit-node=\n force_dns_to_gateway (3 retry)\n SOCKS5
disabled"]
123 SOCKSPATH[" SOCKS5 FALLBACK PATH\n macOS SOCKS5 127.0.0.1:1080\n DNS proxy 127.0.0.1:53\n No exit
node (SKIP_EXIT_NODE=true)"]
124
125 DAEMONS["Start background daemons\n caffeinate -i (sleep prevention)\n DNS Watcher\n WARP Watcher\n(
all write output.log)"]
126
127 RULECOMPILE["Rule compiler status\n(log rule/IP counts)"]
128 WRITEMARKER["write_gateway_active_marker\n(/tmp/nullexit-gateway-active.marker)"]
129 DONE([" Gateway UP"])
130
131 S1 --> S2 --> S3 --> S4 --> S5 --> S6 --> S7
132 S7 --> TSREADY
133 TSREADY -->|"NO (40 NoState)"| ABORT
134 TSREADY -->|"YES"| S8
135 S8 --> S9 --> S10
136 S10 --> PF1 --> PF2 --> PF3
137 PF1 & PF2 & PF3 --> ALLPASS
138 ALLPASS -->|"YES"| EXITPATH
139 ALLPASS -->|"NO"| SOCKSPATH
140 EXITPATH --> DAEMONS
141 SOCKSPATH --> DAEMONS
142 DAEMONS --> RULECOMPILE --> WRITEMARKER --> DONE
143
144 style ABORT fill:#5c1010,color:#fff
145 style EXITPATH fill:#0a3d1a,color:#c0ffc0
146 style SOCKSPATH fill:#3d2d00,color:#ffe0a0
147 style DAEMONS fill:#0a1f3d,color:#c0d8ff
148 style DONE fill:#0a3d1a,color:#fff
149
150
151 ---
152
153 ## 3. toggle.sh STOP Flow
154
155
156 mermaid
157 flowchart TD
158 STOP(["./toggle.sh\n(gateway already ON)"])
159
160 ST1["1. tailscale down\n(disconnect from mesh)"]
161 ST2["2. docker compose down -t 5\n(stop all containers)"]
162 ST3["3. Reset DNS 1.1.1.1\n(networksetup on all services)"]
163 ST4["4. stop_local_dns_proxy\n(kill Python UDP proxy)"]
164 ST5["5. stop_pidfile_daemon\n(kill caffeinate, rm PID)"]
165 ST6["6. stop_pidfile_daemon\n(kill DNS Watcher, rm PID)"]
166 ST7["7. stop_warp_watcher\n(kill WARP Watcher, rm PID)"]
167 ST8["8. clear_gateway_active_marker\n(rm /tmp/nullexit-gateway-active.marker)"]
168 ST9["9. cleanup_network_state\n disable_all_proxies\n Flush DNS cache (dscacheutil)\n Flush stale
routes\n Power-cycle Wi-Fi (offon)"]
169
170 DONE([" Gateway DOWN\nInternet restored"])
171
172 STOP --> ST1 --> ST2 --> ST3 --> ST4
173 ST4 --> ST5 --> ST6 --> ST7 --> ST8 --> ST9 --> DONE
174
175 style DONE fill:#0a3d1a,color:#fff
176
177 ---
178
179 ## 4. Monitoring Layer The 3 Background Daemons
180
181
182 mermaid
183 graph LR
184 subgraph DAEMONS["Background Daemons (all start with gateway, stop on teardown)"]
direction TB

```



```

185     subgraph D1[" Sleep Prevention"]
186         CAFF["caffeinate -i\nPID: /tmp/nullexit-caffeinate.pid\n\nKeeps Mac awake while\ngateway is
187         active"]
188     end
189
190     subgraph D2[" DNS Watcher"]
191         DNSW["Polls networksetup every ~30s\nPID: /tmp/nullexit-dns-watcher.pid\n\nIf DNS drifts from
192         gateway IP\n re-hijacks automatically\n output.log"]
193     end
194
195     subgraph D3[" WARP Watcher"]
196         WARPW["Polls every 30s via:\ndocker compose exec warp wget cdn-cgi/trace\nPID: /tmp/nullexit-
197         warp-watcher.pid\n\nCounts consecutive warp=off\n WARP_FAIL_THRESHOLD (default 6=30s)\n runs recover
198         .sh (nuclear)\n output.log"]
199     end
200
201     end
202
203     subgraph BLIND["What each monitor can/can't see"]
204         B1["WARP Watcher sees:\n Container WARP tunnel state\n Host NIC traffic\n Flash leaks < 5s"]
205         B2["Host-NIC blind spot covered by:\n PF kill-switch (kernel, fail-closed)\n on-demand: diagnose-
206         host-leak.sh\n(--watch to monitor) / sweep.sh"]
207     end
208
209     D3 -->|"covers"| B1
210     B1 -->|"gap closed by"| B2
211
212     style D1 fill:#1a2d0a
213     style D2 fill:#0a1a2d
214     style D3 fill:#2d0a2d
215     style BLIND fill:#1a1a1a,color:#aaa
216
217 ---
218
219 ## 5. Traffic Flow Data Path (Normal Operation)
220
221 mermaid
222 sequenceDiagram
223     participant PHONE as Phone<br/>(Tailnet client)
224     participant TS_HOST as tailscaled<br/>(macOS host)
225     participant TS_CONTAINER as tailscale<br/>(container)
226     participant GLUETUN as Gluetun/warp<br/>(container)
227     participant CF as Cloudflare<br/>WARP Edge
228     participant INTERNET as Internet
229
230     Note over PHONE,INTERNET: Normal EXIT NODE path (all 3 pre-flights passed)
231
232     PHONE->>TS_HOST: WireGuard packet<br/>(encrypted, UDP 41641)
233     TS_HOST->>TS_CONTAINER: route via tun* tailscale0
234     TS_CONTAINER->>GLUETUN: FORWARD chain<br/>(iptables MASQUERADE on tun0)
235     GLUETUN->>TS_HOST: WireGuard tunnel (tun0) egresses macOS Host
236     Note over TS_HOST: macOS pf Firewall applies<br/>Kill-Switch & TCP MSS Clamping
237     TS_HOST->>CF: re-encrypted UDP packet (MSS 1160)
238     CF->>INTERNET: Plain HTTPS from<br/>Cloudflare IP
239
240     Note over PHONE,INTERNET: DNS Resolution path
241     PHONE->>TS_HOST: DNS query
242     TS_HOST->>TS_CONTAINER: UDP:53 TCP:5354 AdGuard :5335
243     TS_CONTAINER->>GLUETUN: AdGuard upstream DNS via tun0
244     GLUETUN->>CF: Encrypted DNS query
245     CF-->>GLUETUN: DNS response
246     GLUETUN-->>TS_CONTAINER: response
247     TS_CONTAINER-->>TS_HOST: response
248     TS_HOST-->>PHONE: DNS answer
249
250 ---
251
252 ## 6. recover.sh Decision Tree
253
254 mermaid
255 flowchart TD
256     INVOKE(["recover.sh called"])
257     MODE{"--post-wake\nflag?"}
258
259     subgraph POSTWAKE[" --post-wake (light refresh, gateway stays UP)"]
260         PW1["Re-hijack DNS gateway IP"]
261         PW2["Flush DNS cache"]
262         PW3["Refresh tailscale exit-node"]
263     end

```

```

261     PW4["Force-recreate warp container\n(if Gluetun healthcheck failed)"]
262     PW5["Leave: containers, Wi-Fi,\ncaffeinate, DNS watcher untouched"]
263 end
264
265 subgraph NUCLEAR[" Default (nuclear gateway torn down)"]
266     N0["Sudo keep-alive loop (background)"]
267     N1["1. tailscale down\n(disconnect from mesh)"]
268     N2["2. Reset DNS 1.1.1.1\n(ALL network services)"]
269     N3["3. disable_all_proxies\n(all interfaces)"]
270     N4["4. Flush DNS cache\n(dscacheutil -flushcache)"]
271     N5["5. Flush stale routes\n(WARP endpoint routes)"]
272     N6a["6a. docker compose down -t 30\n(stop all containers)"]
273     N6b["6b. stop_pidfile_daemon\n(kill caffeinate)"]
274     N6c["6c. stop_pidfile_daemon\n(kill DNS Watcher)"]
275     N7["7. Power-cycle Wi-Fi\n(networksetup offslepp 2on)"]
276     N8["8. Verify internet working\n(curl ifconfig.io)"]
277 end
278
279 DONE_PW([" Gateway refreshed\n(still running)"])
280 DONE_NUC([" Internet restored\n(gateway fully stopped)"])
281
282 INVOKE --> MODE
283 MODE -->|"yes"| PW1
284 PW1 --> PW2 --> PW3 --> PW4 --> PW5 --> DONE_PW
285 MODE -->|"no"| N0
286 N0 --> N1 --> N2 --> N3 --> N4 --> N5
287 N5 --> N6a --> N6b --> N6c --> N7 --> N8 --> DONE_NUC
288
289 style POSTWAKE fill:#0a2d1a,color:#c0ffcc
290 style NUCLEAR fill:#2d0505,color:#ff0c0c
291 style DONE_PW fill:#0a3d1a,color:#fff
292 style DONE_NUC fill:#3d1a00,color:#ffe0cc
293
294
295 ---
296
297 ## 7. Failure Paths & Self-Healing
298
299
300 flowchart LR
301     subgraph EVENTS["Failure / Event"]
302         E1["WARP tunnel drops\n(warp=off)"]
303         E2["Mac wakes from sleep\nor Wi-Fi roams"]
304         E3["DNS drifts from\ngateway IP"]
305         E4["toggle.sh crashes /\nCtrl-C / SIGTERM"]
306         E5["Host-side leak\n(non-tunnel HOST NIC egress)"]
307         E6["Silent NAT timeout\n(ping stops working)"]
308     end
309
310     subgraph DETECTORS["Detector"]
311         D1["WARP Watcher\n(30s poll, docker exec)"]
312         D2["scripts/watcher.sh\n(launchd, always on)"]
313         D3["DNS Watcher\n(30s poll, networksetup)"]
314         D4["cleanup_handler\n(ERR/INT/TERM/HUP trap)"]
315         D5["PF kill-switch\n(kernel, always-on) +\nndiagnose-host-leak.sh --watch\n(on-demand)"]
316         D6["routing-fix.sh\n(30s curl over tun0)"]
317     end
318
319     subgraph ACTIONS["Response"]
320         A1[" threshold consecutive off\n recover.sh (nuclear)"]
321         A2["recover.sh --post-wake\n(gentle refresh)"]
322         A3["Re-hijack DNS\n(networksetup setdnsservers)"]
323         A4["Stop all daemons\nReset DNS 1.1.1.1\nTear down Tailscale"]
324         A5["Non-tunnel egress blocked\nfail-closed at kernel;\n--watch alerts on state change"]
325         A6["Blackhole internet traffic\n+ Drop TUNNEL_FAILED_CLOSED.marker"]
326         A7["Push macOS Desktop Notification\n(via watcher.sh listener)"]
327     end
328
329     E1 --> D1 --> A1
330     E2 --> D2 --> A2
331     E3 --> D3 --> A3
332     E4 --> D4 --> A4
333     E5 --> D5 --> A5
334     E6 --> D6 --> A6
335     A6 -->|"marker detected"| D2
336     D2 --> A7
337
338 style EVENTS fill:#2d1010,color:#ffcccc
339 style DETECTORS fill:#0a1a2d,color:#cce0ff
340 style ACTIONS fill:#0a2d0a,color:#ccffcc
341

```

```

342 ---
343 ---
344
345 ## 8. output.log Who Writes What
346
347 All events converge into a single 'output.log' at the repo root.
348
349 | Writer | Event Types | Format |
350 |-----|-----|-----|
351 | 'toggle.sh' | All startup/shutdown steps, errors, pre-flight results | Plain text with timestamps |
352 | 'recover.sh' | Every recovery step (nuclear or post-wake) | Plain text |
353 | **WARP Watcher** | 'WARP DOWN', 'WARP RECOVERED', 'WARP SHUTDOWN' | '[UTC timestamp]' prefix |
354 | **DNS Watcher** | DNS re-hijack events | Appended inline |
355 | 'docker compose' | Container logs on failure (last 100 lines of warp) | Dumped on ERR path |
356 | 'routing-fix.sh' | (via 'nullexit-logger' inside container) | Structured |
357
358 **Grep cheatsheet:**
359 ``bash
360 grep 'WARP DOWN' output.log # Container-side tunnel drops
361 grep 'WARP SHUTDOWN' output.log # Auto-recovery triggered
362 grep 'WARP RECOVERED' output.log # Self-healed before threshold
363 grep -E 'EXEC|EXIT' output.log # Lifecycle breadcrumbs (last command + exit)
364 ``

```

6 toggle.sh

```

1 #!/bin/bash
2 # toggle.sh nullexit gateway toggle script (macOS + Linux)
3 # Automatically handles Docker containers, Colima, DNS hijacking, and Tailscale exit node routing.
4 #
5 # One file, both OSes: the dispatch below runs the self-contained Linux
6 # implementation (shuf / ss / nmcli / ip / systemd) and exits; on macOS it
7 # falls through to the macOS implementation (jot / netstat / networksetup,
8 # launchd / pf).
9
10 set -e
11
12 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
13
14 #
15 # LINUX TOGGLE self-contained; runs the full toggle then exits before the
16 # macOS body below. Native Linux tooling (shuf / ss / nmcli / ip / systemd).
17 # (Body inlined verbatim at column 0 it contains multiline 'bash -c "'
18 # strings that must NOT be re-indented.)
19 #
20 if [[ "$OSTYPE" != "darwin"* ]]; then
21 # Define log file (used by the lifecycle breadcrumb + run_logged helper)
22 LOG_FILE="$PWD/output.log"
23
24 # --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
25 # TOGGLE_PHASE is updated as the script moves through STOP/START so that if the
26 # process is killed (terminal closed, pkill, OOM) or exits on error, the EXIT
27 # trap records the final exit code AND the phase it died in making an
28 # interrupted START traceable in output.log instead of ending abruptly.
29 TOGGLE_PHASE="init"
30 _log_exit_breadcrumb() {
31     local rc=$?
32     echo "[$(date +%Y-%m-%d %H:%M:%S)] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "
33         $LOG_FILE" 2>/dev/null || true
34 }
35 trap '_log_exit_breadcrumb' EXIT
36
37 # --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
38 # When enabled, mirror a full command-by-command xtrace to output.log. Off by
39 # default. Read directly from .env because common.sh (read_env_var) isn't
40 # sourced yet. Branch on bash version: BASH_XTRACEFD needs >= 4.1 (Linux
41 # usually has bash 5); fall back to teeing stderr on older bash. We never use
42 # the 'exec {var}>>' dynamic-fd form (4.1+ only).
43 DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d '=' -f2- | tr -d "\"" | tr '[:upper:]' '
44     [:lower:]' | tr -d '[:space:]')
45 if [ "$DEBUG_TRACE" = "true" ]; then
46     echo "[$(date +%Y-%m-%d %H:%M:%S)] DEBUG_TRACE enabled xtrace mirrored to output.log (bash ${
47         BASH_VERSION[0]}.${BASH_VERSION[1]})" >> "$LOG_FILE"
48     export PS4='+ [$(SECONDS)s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: ' # subshell-free clock
49     : a $( ) in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
50     if [ "${BASH_VERSION[0]}" -gt 4 ] || { [ "${BASH_VERSION[0]}" -eq 4 ] && [ "${BASH_VERSION[1]}" -ge
51         1 ]; }; then

```

```

47     exec 9>>"$LOG_FILE"
48     export BASH_XTRACEFD=9
49     set -x
50 else
51     # bash < 4.1: plain stderrfile redirect. NOT a process substitution under
52     # set -e, xtrace flowing through a backgrounded 'tee' can abort the script
53     # (failed write/SIGPIPE trips set -e). See devref 15.11.8.
54     exec 2>>"$LOG_FILE"
55     set -x
56 fi
57 fi
58
59 # --- PROCEDURAL PORT GENERATION ---
60 # To prevent static port fingerprinting by malware, we randomize exposed ports
61 # on every boot and persist them to a hidden environment file.
62 PORTS_FILE="/tmp/nullexit-ports.env"
63 if [[ "$1" == "" || "$1" == "--restart" || "$1" == "restart" ]]; then
64
65     # Collision-avoidance function
66     GENERATED_PORTS=""
67     get_free_port() {
68         local port
69         while true; do
70             # Linux uses shuf or $RANDOM instead of jot
71             port=$((shuf -i 10000-65000 -n 1 2>/dev/null || echo $((10000 + RANDOM % 55000))))
72             # 1. Prevent self-collision within this loop
73             if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
74                 continue
75             fi
76             # 2. Prevent collision with ANY process on the machine (root daemons, terminal apps, etc)
77             # We use ss to read the kernel socket table directly.
78             if ! ss -lnt | awk '{print $4}' | grep -q ":$port$"; then
79                 echo "$port"
80                 return
81             fi
82         done
83     }
84
85     export TOR SOCKS_PORT=$(get_free_port)
86     GENERATED_PORTS="$GENERATED_PORTS $TOR SOCKS_PORT"
87
88     export TOR_CONTROL_PORT=$(get_free_port)
89     GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"
90
91     export SOCKS_PROXY_PORT=$(get_free_port)
92     GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"
93
94     export DNS_PROXY_PORT=$(get_free_port)
95
96     echo "export TOR SOCKS_PORT=$TOR SOCKS_PORT" > "$PORTS_FILE"
97     echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
98     echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
99     echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
100 elif [ -f "$PORTS_FILE" ]; then
101     # On STOP, source the existing ports so docker-compose down matches
102     source "$PORTS_FILE"
103 else
104     # Fallbacks if file is lost
105     export TOR SOCKS_PORT=9050
106     export TOR_CONTROL_PORT=9051
107     export SOCKS_PROXY_PORT=1080
108     export DNS_PROXY_PORT=5354
109 fi
110
111 # Start execution timer
112 TOGGLE_START_TIME=$SECONDS
113
114 if [[ "$1" == "restart" ]] || [[ "$1" == "--restart" ]]; then
115     echo "Executing Gateway Restart Sequence..."
116     # We must source common.sh here temporarily to check gateway state before we proceed
117     source "$SCRIPT_DIR/scripts/common.sh"
118     # Use persisted intent via gateway_state (echoes a string, never a return code;
119     # set -e-safe under set -x see devref 15.11.8). On Linux the marker is not
120     # maintained, so this degrades to the is_gateway_active fallback (fast native Docker).
121     if [ "$(gateway_state)" = "running" ]; then
122         echo "Gateway is currently running. Stopping it first..."
123         bash "$0"
124         echo "Gateway stopped. Now starting it..."
125     else
126         echo "Gateway is already stopped. Starting it..."
127     fi

```

```

128     bash "$0"
129     exit 0
130 fi
131
132 # Global flag to track if the script completed successfully
133 SUCCESS_RUN=false
134
135 # Global variable to track the currently running background command PID
136 CURRENT_BG_PID=""
137
138 # Source common bash functions
139 source "$SCRIPT_DIR/scripts/common.sh"
140
141
142 PID_FILE="/tmp/nullexit-cafeinate.pid"
143
144 # Start macOS caffeinate to prevent sleep while gateway is running
145 start_sleep_prevention() {
146     if ! command -v systemd-inhibit >/dev/null 2>&1; then
147         echo " [Warning] systemd-inhibit tool not found. Sleep prevention unavailable."
148         return 1
149     fi
150
151     # Stop any existing sleep prevention process first to avoid duplicates
152     stop_sleep_prevention
153
154     echo " Enabling system sleep prevention (systemd-inhibit)..."
155     # Run systemd-inhibit wrapped in a bash trap to prevent idle system sleep.
156     # Critically, this traps shutdown signals (SIGTERM) and automatically
157     # flushes hijacked DNS back to normal right before the PC powers off.
158     # We do not use recover.sh here because it is a heavy recovery script
159     # which takes too long and gets terminated before it can reset the DNS. Instead,
160     # we surgically and instantly restore normal DNS settings here.
161     nohup bash -c "
162         trap '
163             echo \"[Shutdown Trap] Reverting network settings...\" >> \"\$PWD/output.log\" 2>&1
164             kill \$! 2>/dev/null || true
165             sudo -n resolvectl domain \"\$ACTIVE_SERVICE\" \"\" >> \"\$PWD/output.log\" 2>&1 || true
166             sudo -n resolvectl revert \"\$ACTIVE_SERVICE\" >> \"\$PWD/output.log\" 2>&1 || true
167             resolvectl domain \"\$ACTIVE_SERVICE\" \"\" >> \"\$PWD/output.log\" 2>&1 || true
168             resolvectl revert \"\$ACTIVE_SERVICE\" >> \"\$PWD/output.log\" 2>&1 || true
169             if command -v tailscale >/dev/null 2>&1; then
170                 tailscale up --accept-dns=false --exit-node= >> \"\$PWD/output.log\" 2>&1 &
171                 TS_DOWN_ARGS=\"\"
172                 if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
173                 tailscale down \"\$TS_DOWN_ARGS\" >> \"\$PWD/output.log\" 2>&1 &
174             fi
175             sleep 1
176             echo \"[Shutdown Trap] Cleanup complete.\" >> \"\$PWD/output.log\" 2>&1
177             exit 0
178         ' SIGTERM SIGINT SIGHUP
179         systemd-inhibit --what=sleep --who=nullexit --why=\"gateway running\" --mode=block sleep infinity &
180         wait \$!
181         >> output.log 2>&1 &
182         local caffe_pid=$!
183         echo \"$caffe_pid\" > \"$PID_FILE"
184         echo " Sleep prevention active (PID $caffe_pid). Your PC won't sleep while the gateway is running."
185     }
186
187 # Stop sleep prevention when gateway is stopped
188 stop_sleep_prevention() {
189     if [ -f "$PID_FILE" ]; then
190         local caffe_pid
191         caffe_pid=$(cat "$PID_FILE")
192         if [ -n "$caffe_pid" ] && kill -0 "$caffe_pid" 2>/dev/null && ps -p "$caffe_pid" -o command= 2>/dev/
193             null | grep -q -E "systemd-inhibit|recover.sh"; then
194             echo " Stopping system sleep prevention (systemd-inhibit wrapper PID $caffe_pid)..."
195             kill "$caffe_pid" 2>/dev/null || true
196         fi
197         rm -f "$PID_FILE"
198     fi
199 }
200
201 DNS_WATCHER_PID_FILE="/tmp/nullexit-dns-watcher.pid"
202
203 stop_dns_watcher() {
204     if [ -f "$DNS_WATCHER_PID_FILE" ]; then
205         local watcher_pid
206         watcher_pid=$(cat "$DNS_WATCHER_PID_FILE")
207         if [ -n "$watcher_pid" ] && kill -0 "$watcher_pid" 2>/dev/null; then
208             echo " Stopping background DNS Watcher (PID $watcher_pid)..."

```

```

208     kill -9 "$watcher_pid" 2>/dev/null || true
209 fi
210 rm -f "$DNS_WATCHER_PID_FILE"
211 fi
212 }
213
214 # Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
215 cleanup_handler() {
216     local exit_code=$?
217     local trigger_type="$1"
218     local line_no="$2"
219
220     if [ "$SUCCESS_RUN" = "false" ]; then
221         echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to
222         1.1.1.1..."
223         if [ -n "$ACTIVE_SERVICE" ]; then
224             reset_dns
225         fi
226
227         if [ -n "$CURRENT_BG_PID" ]; then
228             echo "Terminating active command (PID $CURRENT_BG_PID)..."
229             kill -15 "$CURRENT_BG_PID" 2>> output.log || true
230         fi
231
232         # Disconnect host Tailscale and close the GUI application on failure
233         disconnect_tailscale_host
234
235         # Stop local DNS proxy if running
236         stop_local_dns_proxy
237
238         # Stop sleep prevention
239         stop_sleep_prevention
240         stop_dns_watcher
241
242         # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
243         if [ -n "$ACTIVE_SERVICE" ]; then
244             cleanup_network_state
245         fi
246
247         if [ "$trigger_type" = "ERR" ]; then
248             echo -e "\n=====
249             echo "ERROR: Script failed at line $line_no with exit code $exit_code."
250             echo "=====
251             echo "Troubleshooting steps:"
252             echo "1. Run 'colima status' to verify the VM is active."
253             echo "2. Run 'docker compose ps' to check container health."
254             echo "3. Run 'docker compose logs' to view application logs."
255             echo "4. Run 'tailscale status' to check host Tailscale status."
256             echo "=====
257         fi
258     fi
259 }
260
261 trap 'cleanup_handler ERR $LINENO' ERR
262 trap 'cleanup_handler INT' INT
263 trap 'cleanup_handler TERM' TERM
264 trap 'cleanup_handler HUP' HUP
265
266 # NOPASSWD sudo
267 # The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
268 # needed (resolvectl, ip, systemctl, systemd-inhibit, kill, and
269 # python3 scripts/dns-proxy.py for the fallback DNS proxy).
270 # All privileged calls below use 'sudo -n' which will run without prompting.
271 # No credential caching loop is needed.
272
273 # Add Homebrew and standard paths
274 export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
275
276 # Check for local DNS proxy tools (used when tailnet data plane is unavailable)
277 # Uses a Python DNS proxy (handles TCP wire format correctly).
278 # Python3 is built into macOS no external dependencies.
279 DNS_PROXY_BIN=""
280 if command -v python3 >/dev/null 2>&1; then
281     DNS_PROXY_BIN="python3 -I"
282 elif command -v python >/dev/null 2>&1; then
283     DNS_PROXY_BIN="python -I"
284 fi
285
286 # Path to the DNS proxy script (sibling of this script)
287 DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/dns-proxy.py"

```



```

288
289 # Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
290 # ('brew install tailscale' + 'systemctl start tailscaled') NOT the
291 # Tailscale.app GUI so there is no .app bundle or Network Extension sandbox
292 # to launch, no menu-bar click to perform, and no scutil Network Service to
293 # start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
294 # if you ran the install with sudo) and the control socket is always available.
295 TS_BIN=""
296 if command -v tailscale >> output.log 2>&1; then
297     TS_BIN="tailscale"
298 fi
299
300 # Baseline: disable Tailscale DNS management at the daemon level.
301 # 'tailscale set' writes a persistent preference. However, 'tailscale up --reset'
302 # (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
303 # re-enable '--accept-dns=true' and nullify this 'set'.
304 # Therefore, EVERY 'tailscale up --reset' call in this file MUST also pass
305 # '--accept-dns=false' explicitly. See steps 7 and 9 for the fix.
306 #
307 # This initial 'set' is still useful as a baseline for non-reset operations
308 # (e.g. if the user runs 'tailscale up' manually without --reset) and covers
309 # the brief window between this line and the first '--reset' call.
310 #
311 # Runs once at script start while tailscaled is still connected (if it was
312 # left running from a previous toggle), so the pref takes effect before any
313 # disconnection or reconnection logic.
314 if [ -n "$TS_BIN" ]; then
315     $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
316 fi
317
318 # Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
319 # *entire* teardown no need to mess with system network managers.
320 # tailscaled keeps running (as a systemd service), which is intentional:
321 # the next 'tailscale up' then reconnects in seconds rather than waiting for a
322 # slow daemon launch. The 'tailscale down' call also retains the user's
323 # auth state, so we don't force a browser re-login every cycle.
324 #
325 # Defined early (before the if/else branches and referenced by cleanup_handler's
326 # ERR trap) so that any failure before the main logic can still tear down Tailscale.
327 restart_tailscaled_daemon() {
328     echo " [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
329     echo " [Tailscale Recovery] Attempting to restart tailscaled service via systemctl..."
330     if run_with_timeout 15 sudo -n systemctl restart tailscaled >> output.log 2>&1; then
331         echo " [Tailscale Recovery] Successfully restarted tailscaled service."
332     else
333         echo " [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed"
334         echo " ."
335     fi
336     sleep 3
337 }
338 disconnect_tailscale_host() {
339     if [ -n "$TS_BIN" ]; then
340         echo "Disconnecting host Tailscale from mesh (tailscaled stays running as system service)..."
341         local ts_args=""
342         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
343         if ! run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1; then
344             restart_tailscaled_daemon
345             run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1 || true
346         fi
347     fi
348 }
349 ACTIVE_SERVICE=$(get_active_service)
350
351 # Helper to nuke leftover network state after teardown (proxy settings, routing
352 # table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
353 # problem where stale state kills internet after the gateway stops.
354 cleanup_network_state() {
355     echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."
356
357     # 1. Disable any leftover proxy settings (needs root on many macOS versions)
358     # (Skipped for Linux since we don't set global SOCKS proxy)
359     echo " Proxies disabled."
360
361     # 2. Flush DNS cache
362     if command -v resolvectl >> output.log 2>&1; then
363         sudo -n resolvectl flush-caches 2>> output.log || true
364     elif command -v systemd-resolve >> output.log 2>&1; then
365         sudo -n systemd-resolve --flush-caches 2>> output.log || true
366     fi
367     echo " DNS cache flushed."

```

```

368
369 # 3. Flush stale routing table entries
370 if command -v ip >> output.log 2>&1; then
371     sudo -n ip route flush cache >> output.log 2>&1 || true
372 fi
373 echo " Routing table flushed."
374
375 # 4. Power-cycle Wi-Fi to clear any lingering interface state
376 echo " Restarting network interface..."
377 sudo -n ip link set dev "$ACTIVE_SERVICE" down 2>> output.log || true
378 sleep 1
379 sudo -n ip link set dev "$ACTIVE_SERVICE" up 2>> output.log || true
380 echo " Interface power-cycled ($ACTIVE_SERVICE)."
381 sleep 3
382
383 echo "Network state cleanup complete."
384 }
385
386 # Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
387 # during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
388 # 'ts.net' search domain from a previous 'tailscale up --accept-dns=true' run, otherwise macOS
389 # would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
390 # Capture current DNS before resetting so we can check if it was hijacked
391 INITIAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' |
392     head -1 || true)
393
394 echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
395 reset_dns
396
397 # Local Network Surveillance Check
398 echo -e "\n"
399 echo "Local Network Surveillance Check"
400 echo ""
401 # Count populated ARP cache entries to detect network scanning or high congestion
402 if command -v ip >/dev/null 2>&1; then
403     ARP_COUNT=$(ip neigh | grep -iv 'incomplete' | wc -l | tr -d ' ')
404 else
405     # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
406     ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
407 fi
408 if [ "$ARP_COUNT" -gt 15 ]; then
409     echo -e " \033[1;33m[] WARNING: High ARP activity detected ($ARP_COUNT devices in cache).\033[0m"
410     echo -e " \033[1;33m[] This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-
411 scan).\033[0m"
412     echo -e " [] Your traffic remains fully encrypted and invisible.\n"
413 else
414     echo -e " [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
415 fi
416
417 echo "Checking Gateway Status..."
418 # Decide STOP vs START from persisted intent via gateway_state, which echoes a
419 # string (never a return code) to stay set -e-safe under set -x. See devref 15.11.8.
420 # On Linux this falls back to is_gateway_active (fast native Docker).
421 if [ "$(gateway_state)" = "running" ]; then
422     TOGGLE_PHASE="stop"
423     echo -e "\n=====
424     echo "Gateway is RUNNING. Stopping it now..."
425     echo -e "=====
426
427 # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
428 disconnect_tailscale_host
429 echo ""
430
431 echo "Stopping Docker containers..."
432 # '|| true': a disable must always complete even if the Docker daemon is down.
433 run_logged docker compose down -t 5 || true
434
435 echo -e "\nDocker daemon handles container teardown automatically."
436
437 # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
438 # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
439 # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
440 echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)..."
441 reset_dns
442
443 # 3. Stop local DNS proxy if running
444 stop_local_dns_proxy
445
446 # Stop sleep prevention
447 stop_sleep_prevention

```

```

447 stop_dns_watcher
448
449 # 4. Nuke leftover network state so internet actually works after teardown
450 cleanup_network_state
451
452 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
453 echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
454 else
455 TOGGLE_PHASE="start"
456 echo -e "\n=====
457 echo "Gateway is STOPPED. Starting it now..."
458 echo -e "=====
459
460 # 1. Prevent host exit-node deadlock during VM / Container startup
461 disconnect_tailscale_host
462
463 echo -e "\nUsing native Linux Docker..."
464
465 # 4. Clean up corrupted AdGuardHome configurations
466 if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
467     rm -f "adguard/conf/AdGuardHome.yaml"
468     echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
469 fi
470
471
472
473 # 5. Start compose services
474 write_host_ips
475 # Procedurally generated ports have already been exported at the top of the script.
476 echo " [Security] Procedurally randomized proxy ports generated to evade fingerprinting."
477
478 # --- HONEY-PORT TRIPWIRE ---
479 # Generate a random trap port. If any malware does a sequential port scan on localhost,
480 # it will inevitably hit this port. When it does, we instantly fire a desktop alert.
481 # Reap the honey-port from any previous toggle before arming a new one, so the
482 # 'nc -l' tripwire never accumulates one idle listener per toggle-on (15.12.20).
483 pkill -f "nc -l .*127.0.0.1" 2>/dev/null || true
484 pkill -f "nc -l localhost" 2>/dev/null || true
485 HONEY_PORT=$(get_free_port)
486 (
487     if nc -l -p "$HONEY_PORT" 127.0.0.1 >/dev/null 2>&1 || nc -l localhost "$HONEY_PORT" >/dev/null 2>&1;
488     then
489         echo "[$(date +%Y-%m-%d %H:%M:%S)] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown
489         application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$PWD/output.log"
490         if command -v notify-send >/dev/null 2>&1; then
491             notify-send -u critical "nullexit OPSEC Alert" "An unknown application is aggressively scanning
491             your local ports. Malware port-scan suspected."
492         fi
493     fi
494 ) &
495 disown %1 2>/dev/null || true
496 echo " [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."
497
498 echo -e "\nStarting Docker containers..."
499 run_logged docker compose up -d
500
501 # Clean up the one-off rule compiler container so it doesn't clutter the Docker UI
502 docker compose rm -s -f rule-compiler >> output.log 2>&1
503
504 # 6. Wait for the gateway container's Tailscale connection to be ready
505 BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
506 USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')
507
508 echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
509 echo " (This can take 30-60s Tailscale goes through: NoState Starting Running Online)"
510 connected=false
511 consecutive_failures=0
512 for i in {1..60}; do
513     # Run tailscale status and capture its output so the user can see what's happening.
514     # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
515     ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)
516
517     if echo "$ts_output" | grep -q "offers exit node"; then
518         connected=true
519         break
520     fi
521
522     # Track consecutive failures to detect a stuck state.
523     # If we see 40+ consecutive 'NoState' responses, give up early
524     # the daemon is alive but can't connect to the coordination server.
525     if echo "$ts_output" | grep -q "NoState"; then

```

```

525     consecutive_failures=$((consecutive_failures + 1))
526     if [ "$consecutive_failures" -ge 40 ]; then
527         echo ""
528         echo "    [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
529         echo "    Tailscale daemon is running but can't reach the coordination server."
530         break
531     fi
532 else
533     consecutive_failures=0
534 fi
535
536 # Print a condensed status every 10 seconds so the user can see progress
537 if [ $((i % 10)) -eq 0 ]; then
538     ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
539     echo "    [attempt $i/60] $ts_short"
540 elif [ $((i % 5)) -eq 0 ]; then
541     echo -n "[$i]"
542 else
543     echo -n "."
544 fi
545 sleep 1
546 done
547 echo ""
548
549 if [ "$connected" = "true" ]; then
550     echo "Gateway container is online on the Tailscale mesh."
551 elif [ "$BYPASS_PING" = "true" ]; then
552     echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)...
553 "
554 else
555     echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
556     echo "    The container was seen in the following states during the wait:" >&2
557     echo "$ts_output" | sed 's/^/    /' >&2
558     echo "" >&2
559     echo "    This could mean:" >&2
560     echo "    1. Tailscale auth key has expired check TS_AUTHKEY in .env" >&2
561     echo "    2. Network issue inside the container check: docker compose logs tailscale" >&2
562     echo "    3. gluetun VPN tunnel not connecting check: docker compose logs warp | tail -20" >&2
563     echo "    Diagnose manually:" >&2
564     echo "        docker compose exec -T tailscale tailscale status" >&2
565     echo "        docker compose exec -T tailscale tailscale netcheck" >&2
566     cleanup_handler ERR $LINENO
567     exit 1
568 fi
569
570 # 6b. Resolve the gateway's Tailscale IP.
571 # Dynamically query the container for the IP (handles IP changes after re-auth)
572 # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
573 #
574 # CRITICAL: 'docker compose exec -T' on macOS/Colima injects carriage
575 # returns (\r) into captured output. If $TS_IP ends with \r, then
576 # 'networksetup -setdnsservers' silently rejects the malformed IP and
577 # DNS stays at 1.1.1.1 the primary bug this project was hitting.
578 # 'tr -d '\r'' strips the CR; 'awk 'NR==1{print $1; exit}'' takes only
579 # the first field of the first line.
580 TS_IP=""
581 if [ "$connected" = "true" ]; then
582     TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
583     if [ -n "$TS_IP" ]; then
584         echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
585         echo "$TS_IP" > .gateway_ip
586     else
587         echo "    [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
588         TS_IP=$(read_adguard_ip || true)
589     fi
590 else
591     TS_IP=$(read_adguard_ip || true)
592 fi
593
594 # 7. Connect host Mac to Tailscale mesh.
595 # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
596 # system LaunchDaemon via 'systemctl start tailscaled'), the previous five-
597 # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
598 # no menu bar item to click, and no Accessibility permission required.
599 #
600 # CRITICAL: If tailscaled is not running, 'tailscale up' will silently fail.
601 # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
602 # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
603 # are impossible without the host on the mesh).

```

```

604 #
605 # We connect in two phases:
606 # Phase A Join mesh WITHOUT exit node (so pre-flight checks can run)
607 # Phase B Set exit node only after pre-flight checks confirm the gateway works
608 HOST_ON_MESH=false
609 SKIP_EXIT_NODE=false
610 if [ -n "$TS_BIN" ]; then
611     echo -n "Verifying tailscaled is reachable"
612     daemon_ready=false
613     for i in {1..10}; do
614         # 'tailscale status' returns exit code 1 when the daemon is alive but
615         # disconnected ("Tailscale is stopped."). We check the daemon socket
616         # directly as a fallback if the socket exists, tailscaled is running
617         # even if the host isn't connected to the mesh yet.
618         if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
619             daemon_ready=true
620             break
621         fi
622         echo -n "."
623         sleep 1
624     done
625     echo ""
626
627     if [ "$daemon_ready" != "true" ]; then
628         echo -e "\033[0;33m[!] tailscaled daemon is not running. Attempting to start it...\033[0m"
629
630         # Try system-wide daemon (with timeout sudo may prompt for password)
631         if [ "$daemon_ready" != "true" ]; then
632             if run_with_timeout 10 sudo systemctl start tailscaled 2>> output.log; then
633                 echo "Started tailscaled as system LaunchDaemon."
634                 for i in {1..15}; do
635                     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
636                         daemon_ready=true
637                         break
638                     fi
639                     sleep 1
640                 done
641             fi
642         fi
643     fi
644
645     if [ "$daemon_ready" = "true" ]; then
646         echo "tailscaled is responsive (running as a system service)."

```

```

685 echo -n " [1/3] Gateway reachable via Tailscale... "
686 ping_ok=false
687 # The host tailscaled needs a few seconds to propagate the new network map and
688 # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
689 for attempt in {1..45}; do
690     if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
691         ping_ok=true
692         break
693     fi
694     sleep 1
695 done
696 if [ "$ping_ok" = "true" ]; then
697     echo "PASS"
698 else
699     echo "FAIL"
700     SKIP_EXIT_NODE=true
701 fi
702
703 # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
704 echo -n " [2/3] AdGuard DNS via localhost... "
705 if command -v dig >/dev/null 2>&1; then
706     if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 >/dev/null 2>&1; then
707         echo "PASS"
708     else
709         echo "FAIL"
710         SKIP_EXIT_NODE=true
711     fi
712 elif command -v nslookup >/dev/null 2>&1; then
713     if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 >/dev/null 2>&1; then
714         echo "PASS"
715     else
716         echo "FAIL"
717         SKIP_EXIT_NODE=true
718     fi
719 else
720     echo "SKIP (dig/nslookup not available)"
721 fi
722
723 # Check 3: Does the WARP container have external internet?
724 echo -n " [3/3] WARP container internet... "
725 if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace &>/dev/
726 null; then
727     echo "PASS"
728 else
729     echo "FAIL"
730     SKIP_EXIT_NODE=true
731 fi
732
733 # Act based on check results
734 if [ "$SKIP_EXIT_NODE" != "true" ]; then
735     echo -e "\n\033[0;33m[!] Pre-flight checks failed EXIT NODE + DNS HIJACK SKIPPED.\033[0m"
736     echo " Host stays on Tailscale mesh but your internet routes through your normal connection."
737     if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node="$TS_IP" --exit-node-allow-lan-
738 access=true; then
739         echo "Exit node enabled."
740     else
741         echo "[Warning] Failed to set exit node."
742         SKIP_EXIT_NODE=true
743     fi
744 fi
745
746 if [ "$SKIP_EXIT_NODE" = "true" ]; then
747     echo -e "\n\033[0;33m[!] Pre-flight checks failed EXIT NODE + DNS HIJACK SKIPPED.\033[0m"
748     echo " Host stays on Tailscale mesh but your internet routes through your normal connection."
749     echo " Troubleshoot with:"
750     echo "     docker compose logs warp | tail -20"
751     echo "     docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/
752 trace"
753     echo ""
754     echo " Falling back to SOCKS5 proxy for traffic routing..."
755 fi
756
757 elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
758     echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
759     SKIP_EXIT_NODE=true
760 fi
761
762 # 8. Apply DNS Hijacking.
763 # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
764 # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
765 # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
766 if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
767     echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."

```



```

763 echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
764 if force_dns_to_gateway "$TS_IP"; then
765     echo "DNS hijack verified: single server = $TS_IP."
766     start_dns_watcher "$TS_IP"
767 else
768     echo "[Warning] DNS hijack could not be verified."
769     echo "    If DNS is broken, run manually:"
770     echo "        networksetup -setdnsservers \"\$ACTIVE_SERVICE\" \"$TS_IP"
771     echo "        networksetup -setsearchdomains \"\$ACTIVE_SERVICE\" ts.net"
772 fi
773 elif [ -n "$TS_IP" ]; then
774     echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
775     echo "    Trying local DNS proxy for ad-blocking..."
776     if start_local_dns_proxy; then
777         echo "    Ad-blocking active via local DNS proxy through AdGuard."
778     else
779         echo "    Local DNS proxy unavailable. DNS stays at 1.1.1.1."
780         echo "    AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
781     fi
782
783     # Enable SOCKS5 proxy for traffic routing through WARP
784     # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
785     echo -e "\n    Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
786     echo "    (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
787     enable_socks_proxy "$ACTIVE_SERVICE"
788
789     # Verify the SOCKS5 proxy works through WARP
790     echo -n "    Verifying SOCKS5 proxy routes through WARP... "
791     if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn
792     -cgi/trace 2>/dev/null | grep -q "warp=on"; then
793         echo "ok (warp=on)."
794         echo "    All TCP traffic now encrypted through Cloudflare WARP tunnel."
795     else
796         echo "FAILED (proxy not routing through WARP)."
797         echo "    Disabling SOCKS5 proxy internet stays on direct connection."
798         disable_socks_proxy
799         echo "    SOCKS5 proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
800     fi
801 else
802     echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
803 fi
804
805 # 9. Verify host connectivity
806 if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
807     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
808         echo "Host is online on the Tailscale mesh."
809     else
810         echo "[Warning] Host is not responding on Tailscale."
811     fi
812 fi
813
814 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
815 echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."
816
817 # Prevent system from going to sleep while gateway is running
818 start_sleep_prevention
819 fi
820
821 SUCCESS_RUN=true
822
823 # Final DNS state summary
824 if [ -n "$ACTIVE_SERVICE" ]; then
825     FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' |
826     tr ' ' '\n' | true)
827     echo -e "\n"
828     echo -e "DNS STATE: $FINAL_DNS"
829     echo -e ""
830 fi
831
832 if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
833     if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
834         echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
835         echo "    The gateway gracefully fell back to yesterday's cached blocklist."
836         echo "    (Run 'docker logs rule-compiler' later to investigate which link died.)"
837     fi
838 fi
839
840 echo -e "\n"
841 read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
842 if [ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]; then

```

```

842     echo "Reversing gateway state..."
843     exec bash "$@"
844 fi
845
846 echo -e "\nYou can close this terminal window now."
847
848 exit 0
849 fi
850
851 #
852 # macOS TOGGLE  jot / netstat / networksetup, launchd / pf.
853 #
854
855 # Enforce Cryptographic Script Integrity
856 if [ -f "scripts/crypto.sh" ]; then
857     if ! bash scripts/crypto.sh --verify; then
858         exit 1
859     fi
860 fi
861
862 # Define log file
863 LOG_FILE="$PWD/output.log"
864
865 # Rotate log file if it exceeds 1GB (1073741824 bytes).
866 # Sized large on purpose: with DEBUG_TRACE the log holds a full command-by-command
867 # trace effectively the entire compilation/warning history of the framework
868 # which is valuable to keep for post-mortems. Raised from 50MB so always-on
869 # tracing doesn't churn that history away prematurely.
870 if [ -f "$LOG_FILE" ]; then
871     log_size=$(stat -f%z "$LOG_FILE" 2>/dev/null || stat -c%s "$LOG_FILE" 2>/dev/null || echo 0)
872     if [ "$log_size" -gt 1073741824 ]; then
873         mv "$LOG_FILE" "${LOG_FILE}.old" 2>/dev/null || rm -f "$LOG_FILE"
874         echo "[$(date +%Y-%m-%d %H:%M:%S)] Log file exceeded 1GB; rotated to output.log.old" > "$LOG_FILE"
875     fi
876 fi
877
878 # --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
879 # TOGGLE_PHASE is updated as the script moves through STARTUP/STOP/START so that
880 # if the process is killed (window closed, kill, OOM) or exits with an error,
881 # the EXIT trap records the final exit code AND the phase it died in. Previously
882 # an interrupted START left output.log ending abruptly after the teardown with no
883 # indication of what failed this makes every exit traceable.
884 TOGGLE_PHASE="init"
885 _log_exit_breadcrumb() {
886     local rc=$?
887     echo "[$(date +%Y-%m-%d %H:%M:%S)] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "
888         $LOG_FILE" 2>/dev/null || true
889 }
890 trap '_log_exit_breadcrumb' EXIT
891
892 # --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
893 # When enabled, mirror a full command-by-command xtrace to output.log for a
894 # complete post-mortem trail (e.g. a race during --restart while Wi-Fi is
895 # re-associating). Off by default. Read directly from .env here because
896 # common.sh (read_env_var) isn't sourced yet.
897 #
898 # BASH_XTRACEFD (send xtrace to its own fd, keeping the terminal clean) needs
899 # bash >= 4.1. macOS ships /bin/bash 3.2, so we branch:
900 #   - bash >= 4.1 (e.g. Linux): dedicate fd 9 to the log and point xtrace there;
901 #   - bash 3.2 (stock macOS): BASH_XTRACEFD is unsupported, so xtrace always goes
902 #     to stderr. We redirect stderr *straight* to the log file* with a plain
903 #     'exec 2>>"$LOG_FILE"' . This deliberately does NOT use a process
904 #     substitution ('2>>(tee)'): under 'set -e', stderr flowing through the
905 #     backgrounded 'tee' proved to abort the script mid-START on 3.2 (a failed
906 #     write / SIGPIPE in a pipeline returned non-zero and tripped 'set -e', with
907 #     $LINENO mis-blaming the enclosing 'fi' see devref 15.11.8-A/E). The
908 #     tradeoff: on 3.2 the terminal no longer shows live progress while tracing
909 #     (everything is in output.log instead). Acceptable for an opt-in debug mode.
910 # We NEVER use the 'exec {var}>>' dynamic-fd form (4.1+ only), which hard-fails on 3.2.
911 DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d '=' -f2- | tr -d '\n' | tr '[:upper:]' '
912     [:lower:]' | tr -d '[:space:]')
913 if [ "$DEBUG_TRACE" = "true" ]; then
914     echo "[$(date +%Y-%m-%d %H:%M:%S)] DEBUG_TRACE enabled xtrace mirrored to output.log (bash ${
915         BASH_VERSION[0]}.${BASH_VERSION[1]})" >> "$LOG_FILE"
916     # Timestamped, location-aware xtrace prompt (works on 3.2 and 4.x).
917     export PS4='+ [${SECONDS}s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: ' # subshell-free clock
918     : a $(()) in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
919     if [ "${BASH_VERSION[0]}" -gt 4 ] || { [ "${BASH_VERSION[0]}" -eq 4 ] && [ "${BASH_VERSION[1]}" -ge
920         1 ]; }; then
921         exec 9>>"$LOG_FILE"

```

```

918     export BASH_XTRACEFD=9
919     set -x
920 else
921     # bash 3.2 (stock macOS): plain stderrfile redirect (no process substitution).
922     exec 2>>"$LOG_FILE"
923     set -x
924 fi
925 fi
926
927 # --- MAIN EXECUTION LOGGING ---
928 echo -e "\n===== " >> "$LOG_FILE"
929 echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh invoked" >> "$LOG_FILE"
930 echo "===== " >> "$LOG_FILE"
931
932 # (Previously this script chmod 600'd .env at start to "unlock" it for docker compose,
933 # and chmod 000'd it on exit. That pattern broke docker compose on macOS/Colima when
934 # the umask produced 0600 the script removed its own ability to read .env before it could even proceed.)
935
936 # --- PROCEDURAL PORT GENERATION ---
937 # To prevent static port fingerprinting by malware, we randomize exposed ports
938 # on every boot and persist them to a hidden environment file.
939 PORTS_FILE="/tmp/nullexit-ports.env"
940 if [[ "$1" == "" || "$1" == "--restart" ]]; then
941
942     # Collision-avoidance function
943     GENERATED_PORTS=""
944     get_free_port() {
945         local port
946         while true; do
947             port=$(jot -r 1 10000 65000)
948             # 1. Prevent self-collision within this loop
949             if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
950                 continue
951             fi
952             # 2. Prevent collision with ANY process on the Mac (root daemons, terminal apps, etc)
953             # We use netstat instead of lsof because netstat reads the kernel socket table
954             # directly and doesn't require sudo to see ports bound by other users/root.
955             if ! netstat -an -p tcp | awk '{print $4}' | grep -q "\.$port$"; then
956                 echo "$port"
957                 return
958             fi
959         done
960     }
961
962     export TOR SOCKS_PORT=$(get_free_port)
963     GENERATED_PORTS="$GENERATED_PORTS $TOR SOCKS_PORT"
964
965     export TOR_CONTROL_PORT=$(get_free_port)
966     GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"
967
968     export SOCKS_PROXY_PORT=$(get_free_port)
969     GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"
970
971     export DNS_PROXY_PORT=$(get_free_port)
972
973     echo "export TOR SOCKS_PORT=$TOR SOCKS_PORT" > "$PORTS_FILE"
974     echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
975     echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
976     echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
977     ADGUARD_PASS=$(grep -E "^ADGUARD_PASSWORD=" .env 2>/dev/null | cut -d'=' -f2- | tr -d "\'\"")
978     echo "export TOR_PASSWORD=${ADGUARD_PASS:-nullexit}" >> "$PORTS_FILE"
979 elif [ -f "$PORTS_FILE" ]; then
980     # On STOP, source the existing ports so docker-compose down matches
981     source "$PORTS_FILE"
982 else
983     # Fallbacks if file is lost
984     export TOR SOCKS_PORT=9050
985     export TOR_CONTROL_PORT=9051
986     export SOCKS_PROXY_PORT=1080
987     export DNS_PROXY_PORT=5354
988 fi
989
990 # Start execution timer
991 TOGGLE_START_TIME=$SECONDS
992
993 if [[ "$1" == "--restart" ]]; then
994     echo "Executing Gateway Restart Sequence..."
995     # We must source common.sh here temporarily to check gateway state before we proceed
996     source "$SCRIPT_DIR/scripts/common.sh"
997     # Use persisted intent (marker) via gateway_state, which echoes a string (never
998     # a return code) to stay set -e-safe under set -x. See devref 15.11.8.

```

```

999     if [ "$(gateway_state)" = "running" ]; then
1000         echo "Gateway is currently running. Stopping it first..."
1001         bash "$0"
1002         echo "Gateway stopped. Now starting it..."
1003     else
1004         echo "Gateway is already stopped. Starting it..."
1005     fi
1006     bash "$0"
1007     exit 0
1008 fi
1009
1010 # Global flag to track if the script completed successfully
1011 SUCCESS_RUN=false
1012
1013 # Global variable to track the currently running background command PID
1014 CURRENT_BG_PID=""
1015
1016 # Source common bash functions
1017 source "$SCRIPT_DIR/scripts/common.sh"
1018
1019 # Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
1020 LOCK_FILE="/tmp/nullexit-toggle.lock"
1021 if [ -f "$LOCK_FILE" ]; then
1022     LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
1023     if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
1024         if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
1025             die "Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running."
1026         fi
1027     fi
1028 fi
1029 echo "$$" > "$LOCK_FILE"
1030
1031 # Helper function to run a GUI command as the logged-in console user
1032 # (Prevents permission failures if the script is run with sudo)
1033 run_gui_cmd() {
1034     local console_user
1035     console_user=$(stat -f '%Su' /dev/console 2>> output.log || echo "$SUDO_USER")
1036     if [ -z "$console_user" ] || [ "$console_user" = "root" ]; then
1037         console_user=$(logname 2>> output.log || echo "$USER")
1038     fi
1039
1040     if [ -n "$console_user" ] && [ "$console_user" != "root" ] && [ "$EUID" -eq 0 ]; then
1041         sudo -u "$console_user" "$@"
1042     else
1043         "$@"
1044     fi
1045 }
1046
1047 # Active-state marker: written at end of START path so external watchers
1048 # (see launchd/com.nullexit.wake-recovery.plist + scripts/watcher.sh +
1049 # devref 10.29) know the gateway is currently up. They only fire
1050 # recover.sh --post-wake when this file is present. Cleared at the top
1051 # of the STOP path and inside cleanup_handler on any error/signal path
1052 # so a half-broken gateway never re-triggers post-wake refreshes.
1053 write_gateway_active_marker() {
1054     date -u +%FT%TZ > /tmp/nullexit-gateway-active.marker
1055     # Set the watcher debounce timestamp to now to prevent a redundant post-startup recovery run.
1056     date +%s > /tmp/nullexit-watcher.last-recovery 2>/dev/null || true
1057 }
1058
1059 clear_gateway_active_marker() {
1060     rm -f /tmp/nullexit-gateway-active.marker
1061     rm -f "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/TUNNEL_FAILED_CLOSED.marker"
1062 }
1063
1064 # Capture whether the gateway-active marker exists BEFORE the defensive clear
1065 # below wipes it. The STOP-vs-START decision (gateway_should_be_running) reads
1066 # this captured value persisted *intent* instead of a live Docker probe, so
1067 # a disable is instant and works even when Colima/Docker is down. Must be read
1068 # before clear_gateway_active_marker or it would always look "stopped".
1069 #
1070 # Written as default-assign + 'if' WITHOUT an 'else': under 'set -e' (with the
1071 # DEBUG_TRACE xtrace redirect active) an 'if [ -f ]; then ; else ; fi' whose
1072 # condition is false was aborting the script here at init on macOS /bin/bash 3.2
1073 # the false '[ -f ]' status leaked through the compound. An 'if' with no 'else'
1074 # yields status 0 when the condition is false, so nothing leaks.
1075 GATEWAY_MARKER_AT_STARTUP="no"
1076 if [ -f "$MARKER_FILE" ]; then
1077     GATEWAY_MARKER_AT_STARTUP="yes"
1078 fi
1079

```

```

1080 # Defensive: clear any stale marker from a prior crashed/aborted run. If a
1081 # previous toggle.sh wrote the marker but was OOM-killed in the middle of the
1082 # START block (before the END-of-START write) or if the user rebooted mid
1083 # session this prevents the watcher from firing post-wake against a
1084 # non-running gateway on every wake event. Placed AFTER the function
1085 # definitions so bash resolves the call at parse time.
1086 clear_gateway_active_marker
1087
1088 PID_FILE="$PID_CAFFEINATE"
1089
1090 # Start macOS caffeinate to prevent sleep while gateway is running
1091 start_sleep_prevention() {
1092     if ! command -v caffeinate >/dev/null 2>&1; then
1093         echo " [Warning] caffeinate tool not found. Sleep prevention unavailable."
1094         return 1
1095     fi
1096
1097     # Stop any existing sleep prevention process first to avoid duplicates
1098     stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
1099
1100     echo " Enabling system sleep prevention (caffeinate)..."
1101     # Run caffeinate wrapped in a bash trap to prevent idle system sleep.
1102     # Critically, this traps macOS shutdown signals (SIGTERM) and automatically
1103     # flushes hijacked DNS back to normal right before the Mac powers off.
1104     # We do not use recover.sh here because it is a heavy recovery script (managing
1105     # containers, restarting tailscaled, etc.) which takes too long and gets
1106     # terminated by macOS before it can reset the DNS. Instead, we surgically and
1107     # instantly restore normal DNS and proxy settings here.
1108     nohup bash -c "
1109         trap '
1110             echo \"[Shutdown Trap] Reverting network settings...\" >> \"\$PWD/output.log\" 2>&1
1111             kill \$! 2>/dev/null || true
1112             sudo -n networksetup -setdnsservers \"\$ACTIVE_SERVICE\" 1.1.1.1 >> \"\$PWD/output.log\" 2>&1 || true
1113             sudo -n networksetup -setdnsservers \"\$ENO_SERVICE\" 1.1.1.1 >> \"\$PWD/output.log\" 2>&1 || true
1114             sudo -n networksetup -setsearchdomains \"\$ACTIVE_SERVICE\" \"Empty\" >> \"\$PWD/output.log\" 2>&1 ||
1115             true
1116             sudo -n networksetup -setsearchdomains \"\$ENO_SERVICE\" \"Empty\" >> \"\$PWD/output.log\" 2>&1 ||
1117             true
1118             sudo -n networksetup -setsocksfirewallproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\"
1119             2>&1 || true
1120             sudo -n networksetup -setsocksfirewallproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1
1121             || true
1122             sudo -n networksetup -setwebproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1123             sudo -n networksetup -setwebproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1124             sudo -n networksetup -setsecurewebproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 ||
1125             true
1126             sudo -n networksetup -setsecurewebproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 ||
1127             true
1128             if command -v tailscale >/dev/null 2>&1; then
1129                 tailscale up >> \"\$PWD/output.log\" 2>&1 || true
1130                 tailscale set --accept-dns=false --exit-node= >> \"\$PWD/output.log\" 2>&1 || true
1131                 TS_DOWN_ARGS=\"\"
1132                 if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
1133                 tailscale down \"\$TS_DOWN_ARGS\" >> \"\$PWD/output.log\" 2>&1 &
1134             fi
1135             sleep 1
1136             echo \"[Shutdown Trap] Cleanup complete.\" >> \"\$PWD/output.log\" 2>&1
1137             exit 0
1138         ' SIGTERM SIGINT SIGHUP
1139         caffeinate -i &
1140         wait \$!
1141         \" >> output.log 2>&1 &
1142         local caffe_pid=$!
1143         echo \"$caffe_pid\" > \"$PID_FILE"
1144         echo " Sleep prevention active (PID $caffe_pid). Your Mac won't sleep while the gateway is running."
1145     }
1146
1147 DNS_WATCHER_PID_FILE="$PID_DNS_WATCHER"
1148 WARP_WATCHER_PID_FILE="/tmp/nullexit-warp-watcher.pid"
1149
1150 # Background WARP tunnel liveness monitor. Polls cdn-cgi/trace every 30s
1151 # while healthy, but accelerates to polling every 5s if a failure is detected.
1152 # Always logs state transitions to output.log. When warp=off persists for
1153 # WARP_FAIL_THRESHOLD consecutive polls (default 6 = 30s of downtime), the watcher
1154 # triggers a nuclear recovery: runs recover.sh to tear down the entire
1155 # gateway disconnecting Tailscale, stopping containers, resetting DNS,
1156 # and power-cycling Wi-Fi. Note: This minimizes exposure window, but

```

```

1155 # only the pf kill switch guarantees absolutely zero leakage on drop.
1156 #
1157 # The threshold (default 6) is statistically chosen: even at an
1158 # unrealistically high 10% false-positive rate per poll, P(6 consecutive
1159 # false positives) = 0.1^6 0.0001%. Real-world false-positive rates
1160 # are far lower, making 30s of continuous downtime overwhelmingly likely
1161 # to be a genuine outage.
1162 #
1163 # Configurable via .env: WARP_FAIL_THRESHOLD=3 (15s, more aggressive)
1164 start_warp_watcher() {
1165     stop_warp_watcher
1166     # Parse threshold from .env (same pattern as GATEWAY_MSS, GATEWAY_BYPASS_PING, etc.)
1167     local threshold
1168     threshold=$(read_env_var WARP_FAIL_THRESHOLD)
1169     if [ -z "$threshold" ]; then
1170         threshold="${WARP_FAIL_THRESHOLD:-6}"
1171     fi
1172     local trigger_file="/tmp/nullexit-warp-shutdown-triggered"
1173     rm -f "$trigger_file"
1174     echo "    Starting background WARP Watcher (polling every 30s [accelerating to 5s on failure], shutdown
        after ${threshold} consecutive failures)..."
1175     nohup bash -c "
1176         trap 'exit 0' SIGTERM SIGINT SIGHUP
1177         last_state='on'
1178         consec_off=0
1179         threshold='$threshold'
1180         trigger_file='$trigger_file'
1181         out_log='$PWD/output.log'
1182         recover_bin='$PWD/recover.sh'
1183         inhibit_file='/tmp/nullexit-warp-inhibit.marker'
1184         while true; do
1185             # Inhibit check
1186             # recover.sh --post-wake writes this marker while force-recreating the
1187             # warp container. During that window, the container is intentionally
1188             # down don't count it as a failure or we'll fire nuclear recover.sh
1189             # and kill a gateway that was in the process of healing itself.
1190             if [ -f "\$inhibit_file" ]; then
1191                 consec_off=0
1192                 sleep 5
1193                 continue
1194             fi
1195
1196             state='off'
1197             if docker compose exec -T warp wget -q0- --timeout=5 \
1198                 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
1199                 | grep -q 'warp=on'; then
1200                 state='on'
1201             fi
1202
1203             # State-transition logging
1204             if [ "\$state" != "\$last_state" ]; then
1205                 if [ "\$state" = 'off' ]; then
1206                     echo "\[$(date -u +%FT%TZ)] WARP DOWN    cdn-cgi/trace reports warp=off\" >> \"\$out_log\"
1207                 fi
1208                 last_state="\$state"
1209             fi
1210
1211             # Consecutive-failure tracking + auto-shutdown
1212             if [ "\$state" = 'off' ]; then
1213                 consec_off=$((consec_off + 1))
1214                 if [ "\$consec_off" -ge "\$threshold" ] && [ ! -f "\$trigger_file" ]; then
1215                     touch "\$trigger_file"
1216                     echo "\[$(date -u +%FT%TZ)] WARP SHUTDOWN    \$consec_off consecutive failures (threshold=\
\$threshold). Running recover.sh to kill the gateway and restore normal internet. Your IP is NO
LONGER Cloudflare.\" >> \"\$out_log\"
1217                     # Nuclear recovery: tear down the entire gateway.
1218                     # recover.sh resets DNS 1.1.1.1, disconnects Tailscale,
1219                     # stops Docker containers, flushes routes, power-cycles Wi-Fi.
1220                     bash "\$recover_bin" --auto >> \"\$out_log\" 2>&1 || true
1221                     echo "\[$(date -u +%FT%TZ)] WARP SHUTDOWN    recover.sh completed, watcher exiting. Your IP is
now your real ISP IP.\" >> \"\$out_log\"
1222                     exit 0
1223                 fi
1224             else
1225                 if [ "\$consec_off" -gt 0 ]; then
1226                     echo "\[$(date -u +%FT%TZ)] WARP RECOVERED    warp=on after \$consec_off consecutive off
readings (threshold was \$threshold)\" >> \"\$out_log\"
1227                 fi
1228                 consec_off=0
1229             fi
1230

```



```

1231     if [ "\$state\" = \"off\" ]; then
1232         sleep 5
1233     else
1234         sleep 30
1235     fi
1236 done
1237 " >> output.log 2>&1 &
1238 echo $! > "$WARP_WATCHER_PID_FILE"
1239 }
1240
1241 stop_warp_watcher() {
1242     if [ -f "$WARP_WATCHER_PID_FILE" ]; then
1243         local wp
1244         wp=$(cat "$WARP_WATCHER_PID_FILE")
1245         if [ -n "$wp" ] && kill -0 "$wp" 2>/dev/null; then
1246             echo "Stopping background WARP Watcher (PID $wp)..."
1247             kill "$wp" 2>/dev/null || true
1248         fi
1249         rm -f "$WARP_WATCHER_PID_FILE"
1250     fi
1251 }
1252
1253
1254
1255
1256
1257 # Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
1258 cleanup_handler() {
1259     local exit_code=$?
1260     local trigger_type="$1"
1261     local line_no="$2"
1262
1263     if [ "$SUCCESS_RUN" = "false" ]; then
1264         echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1..."
1265         if [ -n "$ACTIVE_SERVICE" ]; then
1266             reset_dns
1267         fi
1268
1269         if [ -n "$CURRENT_BG_PID" ]; then
1270             echo "Terminating active command (PID $CURRENT_BG_PID)..."
1271             kill -15 "$CURRENT_BG_PID" 2>> output.log || true
1272         fi
1273
1274         # Disconnect host Tailscale and close the GUI application on failure
1275         disconnect_tailscale_host
1276
1277         # Stop local DNS proxy if running
1278         stop_local_dns_proxy
1279
1280         # Stop sleep prevention
1281         stop_pidfile_daemon "/tmp/nullexit-cafeinate.pid" "system sleep prevention"
1282         stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
1283         stop_warp_watcher
1284         clear_gateway_active_marker
1285
1286         # Capture warp logs on failure for debugging before teardown
1287         if [ "$trigger_type" = "ERR" ]; then
1288             echo -e "\n--- WARP FAILURE LOGS ---" >> output.log
1289             docker compose logs warp --tail=100 >> output.log 2>&1 || true
1290             echo "-----" >> output.log
1291         fi
1292
1293         # Best-effort container teardown on failure
1294         docker compose down --remove-orphans -t 30 2>> output.log || true
1295
1296         # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
1297         if [ -n "$ACTIVE_SERVICE" ]; then
1298             cleanup_network_state
1299         fi
1300
1301
1302         if [ "$trigger_type" = "ERR" ]; then
1303             echo -e "\n=====
1304             echo "ERROR: Script failed at line $line_no with exit code $exit_code."
1305             echo "=====
1306             echo "Troubleshooting steps:"
1307             echo "1. Run 'colima status' to verify the VM is active."
1308             echo "2. Run 'docker compose ps' to check container health."
1309             echo "3. Run 'docker compose logs' to view application logs."
1310             echo "4. Run 'tailscale status' to check host Tailscale status."

```

```

1311     echo "=====
1312 fi
1313 rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
1314 fi
1315 }
1316
1317 trap 'cleanup_handler ERR $LINENO' ERR
1318 trap 'cleanup_handler INT' INT
1319 trap 'cleanup_handler TERM' TERM
1320 trap 'cleanup_handler HUP' HUP
1321
1322 # NOPASSWD sudo
1323 # The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
1324 # needed (networksetup, pfctl, route, ifconfig, dscacheutil, killall, pkill,
1325 # and python3 -I scripts/dns-proxy.py for the fallback DNS proxy).
1326 # All privileged calls below use 'sudo -n' which will run without prompting.
1327 # No credential caching loop is needed.
1328
1329 # Add Homebrew and standard paths
1330 setup_standard_path
1331
1332 # Check for local DNS proxy tools (used when tailnet data plane is unavailable)
1333 # Uses a Python DNS proxy (handles TCP wire format correctly).
1334 # Python3 is built into macOS no external dependencies.
1335 DNS_PROXY_BIN=""
1336 if command -v python3 >> output.log 2>&1; then
1337     DNS_PROXY_BIN="python3 -I"
1338 elif command -v python >> output.log 2>&1; then
1339     DNS_PROXY_BIN="python -I"
1340 fi
1341
1342 # Path to the DNS proxy script (sibling of this script)
1343 DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/dns-proxy.py"
1344
1345 # Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
1346 # ('brew install tailscale' + 'brew services start tailscale') NOT the
1347 # Tailscale.app GUI so there is no .app bundle or Network Extension sandbox
1348 # to launch, no menu-bar click to perform, and no scutil Network Service to
1349 # start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
1350 # if you ran the install with sudo) and the control socket is always available.
1351 TS_BIN=""
1352 if command -v tailscale >> output.log 2>&1; then
1353     TS_BIN="tailscale"
1354 fi
1355
1356 # Baseline: disable Tailscale DNS management at the daemon level.
1357 # 'tailscale set' writes a persistent preference. However, 'tailscale up --reset'
1358 # (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
1359 # re-enable '--accept-dns=true' and nullify this 'set'.
1360 # Therefore, EVERY 'tailscale up --reset' call in this file MUST also pass
1361 # '--accept-dns=false' explicitly. See steps 7 and 9 for the fix.
1362 #
1363 # This initial 'set' is still useful as a baseline for non-reset operations
1364 # (e.g. if the user runs 'tailscale up' manually without --reset) and covers
1365 # the brief window between this line and the first '--reset' call.
1366 #
1367 # Runs once at script start while tailscaled is still connected (if it was
1368 # left running from a previous toggle), so the pref takes effect before any
1369 # disconnection or reconnection logic.
1370 if [ -n "$TS_BIN" ]; then
1371     $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
1372 fi
1373
1374 # Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
1375 # *entire* teardown no scutil tunnel stop, no Tailscale.app GUI quit, no pkill.
1376 # tailscaled keeps running (as a LaunchAgent / LaunchDaemon), which is intentional:
1377 # the next 'tailscale up' then reconnects in seconds rather than waiting for a
1378 # slow .app launch + Network Extension activation. The 'tailscale down' call also
1379 # retains the user's auth state, so we don't force a browser re-login every cycle.
1380 #
1381 # Defined early (before the if/else branches and referenced by cleanup_handler's
1382 # ERR trap) so that any failure before the main logic can still tear down Tailscale.
1383
1384
1385 # Wait for Network Readiness
1386 # On --restart, the gateway tears down while the host network interface may
1387 # still be reconnecting (Wi-Fi, Ethernet, USB tethering, etc.). If we proceed
1388 # too early, get_active_service returns nothing, routing targets the wrong
1389 # interface, and DNS/WARP setup silently fails.
1390 #
1391 # We use 'route get default' to detect a valid default gateway rather than

```

```

1392 # checking a specific interface (e.g. en0) this generalizes across all
1393 # connection types: Wi-Fi, Ethernet, USB-C adapters, iPhone tethering, etc.
1394 # The moment the OS has any routable default gateway, we proceed.
1395 _net_wait_secs=0
1396 _net_timeout=60
1397 echo "Waiting for network connectivity before starting..."
1398 while true; do
1399     if route get default 2>/dev/null | grep -q 'gateway:~'; then
1400         _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
1401         echo "    Network ready (default gateway: $_gw)."
1402         break
1403     fi
1404     if [ "$_net_wait_secs" -ge "$_net_timeout" ]; then
1405         echo "    [Warning] No default gateway after ${_net_timeout}s proceeding anyway."
1406         break
1407     fi
1408     sleep 2
1409     _net_wait_secs=$(( _net_wait_secs + 2 ))
1410 done
1411
1412 ACTIVE_SERVICE=$(get_active_service)
1413
1414 # Resolve the service name for en0 (usually "Wi-Fi") macOS scutil DNS resolver
1415 # is commonly scoped to en0, so per-service DNS changes on other interfaces are
1416 # ignored unless en0's service is also updated.
1417 ENO_SERVICE=$(get_en0_service)
1418
1419 # Helper to add host-side bypass routes for the WARP WireGuard endpoints.
1420 # This prevents an infinite recursive routing loop (tunnel loop) where
1421 # the container's WireGuard packets exit the VM, reach the host, match the
1422 # default route (the exit node), and get routed back into the tunnel.
1423 # We route via the gateway IP (not interface) to ensure packets are routed
1424 # properly even when the host's default route changes to the Tailscale interface.
1425
1426 # Helper to manually override the host's default route to point through the
1427 # Tailscale utun* interface. Standalone tailscaled on macOS (Homebrew version)
1428 # lacks the platform-specific integration to override the default gateway
1429 # automatically when --exit-node is used.
1430 setup_exit_node_routing() {
1431     local ts_iface
1432     ts_iface=$(ifconfig | awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':')
1433
1434     if [ -n "$ts_iface" ]; then
1435         echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
1436         local host_mtu
1437         host_mtu=$(read_env_var HOST_MTU)
1438         host_mtu=${host_mtu:-1200}
1439
1440         # Lower the host Tailscale MTU (defaults to 1200) to prevent fragmentation when passing
1441         # through the container's WARP tunnel (which also has an MTU of 1280).
1442         sudo -n ifconfig "$ts_iface" mtu "$host_mtu" >> output.log 2>&1 || true
1443
1444         # DO NOT delete the physical default route! It crashes Colima.
1445         # Use the /1 trick to mathematically override the default route.
1446         sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
1447         sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
1448         sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
1449         sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
1450
1451         if netstat -nr | grep -E -q "^(0\.\0\.\0\.\0|1|0/1)[[:space:]]+.*$ts_iface"; then
1452             echo "    [] Default route successfully overridden via $ts_iface."
1453         else
1454             echo "    [!] Failed to verify exit node routing on $ts_iface."
1455         fi
1456     else
1457         echo "[Warning] Could not detect Tailscale utun interface for host routing."
1458     fi
1459 }
1460
1461 # Helper to nuke leftover network state after teardown (proxy settings, routing
1462 # table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
1463 # problem where stale state kills internet after the gateway stops.
1464 cleanup_network_state() {
1465     echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."
1466
1467     # 1. Disable any leftover proxy settings (needs root on many macOS versions)
1468     for svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
1469         disable_all_proxies "$svc"
1470     done
1471     echo "    Proxies disabled."
1472

```

```

1473 # 2. Flush DNS cache
1474 if command -v dscacheutil >> output.log 2>&1; then
1475     sudo -n dscacheutil -flushcache 2>> output.log || true
1476 fi
1477 sudo -n killall -HUP mDNSResponder 2>> output.log || true
1478 echo "    DNS cache flushed."
1479
1480 # 3. Flush stale routing table entries
1481 if command -v route >> output.log 2>&1; then
1482     sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
1483     sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
1484 fi
1485 remove_warp_bypass_routes
1486 remove_apple_split_routes
1487 disable_killswitch
1488 echo "    Routing table and firewall flushed."
1489
1490 # 4. Power-cycle Wi-Fi to clear any lingering interface state
1491 WIFI_PORT=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline;
1492     print $2}')
1493 if [ -n "$WIFI_PORT" ]; then
1494     sudo -n networksetup -setairportpower "$WIFI_PORT" off 2>> output.log || true
1495     sleep 2
1496     sudo -n networksetup -setairportpower "$WIFI_PORT" on 2>> output.log || true
1497     echo "    Wi-Fi power-cycled (interface $WIFI_PORT)."
1498     sleep 3
1499 else
1500     echo "    Wi-Fi interface not detected; bouncing en0..."
1501     sudo -n ifconfig en0 down 2>> output.log || true
1502     sudo -n ifconfig en0 up 2>> output.log || true
1503 fi
1504 # Force restart sharing services to prevent AirDrop / AirPlay discovery freezes
1505 # after network interface/DNS changes.
1506 echo "    Resetting macOS sharing services (AirDrop/AirPlay)..."
1507 reset_sharing_services
1508
1509 echo "Network state cleanup complete."
1510 }
1511
1512 SOCKS_PROXY_PORT=${SOCKS_PROXY_PORT}
1513
1514 # Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
1515 # during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
1516 # 'ts.net' search domain from a previous 'tailscale up --accept-dns=true' run, otherwise macOS
1517 # would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
1518
1519
1520 echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
1521 reset_dns
1522
1523
1524
1525
1526 echo "Checking Gateway Status..."
1527 HIJACK_HOST=$(read_env_var GATEWAY_HIJACK_HOST | tr '[:upper:]' '[:lower:]')
1528
1529 # Decide STOP vs START from persisted intent (marker), not a live Docker probe.
1530 # This makes *disable* instant and robust even when Colima/Docker is down.
1531 #
1532 # CRITICAL (set -e + set -x on bash 3.2): gateway_state ECHOES "running"/"stopped"
1533 # and always returns 0, so we branch on the STRING never on a function's exit
1534 # code. A helper that 'return 1's trips a spurious 'set -e' abort under 'set -x'
1535 # (DEBUG_TRACE=true) even when guarded, hence the echo contract. See devref 15.11.8.
1536 if [ "$(gateway_state)" = "running" ]; then
1537     TOGGLE_PHASE="stop"
1538     echo -e "\n=====
1539     echo "Gateway is RUNNING. Stopping it now..."
1540     echo -e "=====
1541
1542 # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
1543 if [[ "$HIJACK_HOST" != "false" ]]; then
1544     disconnect_tailscale_host
1545     echo ""
1546 fi
1547
1548 echo "Stopping Docker containers..."
1549 # '|| true': a disable must always complete its teardown even if the Docker
1550 # daemon / Colima VM is already down (compose down would exit non-zero and,
1551 # under set -e, abort the STOP). The START path intentionally does NOT guard
1552 # its 'docker compose up' a failed start SHOULD trigger cleanup.

```

```

1553 run_logged docker compose down --remove-orphans -t 30 || true
1554
1555 # Only stop Colima if the user explicitly opted in (false by default) to prevent breaking their other
1556 # Docker dev projects
1557 STOP_COLIMA=$(read_env_var STOP_COLIMA_ON_EXIT | tr '[:upper:]' '[:lower:]')
1558 if [[ "$STOP_COLIMA" == "true" ]]; then
1559     echo -e "\nStopping Colima VM to free up host RAM and battery..."
1560     run_with_timeout 30 colima stop >> output.log 2>&1 || echo "Warning: Failed to stop Colima gracefully"
1561 else
1562     echo -e "\nLeaving Colima running. (Set STOP_COLIMA_ON_EXIT=true in .env to change this)"
1563 fi
1564
1565 if [[ "$HIJACK_HOST" != "false" ]]; then
1566     # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
1567     # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
1568     # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
1569     echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)..."
1570     reset_dns
1571
1572     # 3. Stop local DNS proxy if running
1573     stop_local_dns_proxy
1574
1575     # Stop background daemons
1576     stop_pidfile_daemon "/tmp/nullexit-cafeinate.pid" "system sleep prevention"
1577     stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
1578     stop_warp_watcher
1579     clear_gateway_active_marker
1580
1581     # 4. Nuke leftover network state so internet actually works after teardown
1582     cleanup_network_state
1583 fi
1584
1585 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
1586 echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
1587 else
1588     TOGGLE_PHASE="start"
1589     START_GATEWAY=true
1590     echo "[$(date +%Y-%m-%d %H:%M:%S')] Action: STARTUP GATEWAY" >> "$LOG_FILE"
1591     echo -e "\n=====
1592     echo "Gateway is STOPPED. Starting it now..."
1593     echo -e "=====
1594
1595     # 1. Prevent host exit-node deadlock during VM / Container startup
1596     if [[ "$HIJACK_HOST" != "false" ]]; then
1597         disconnect_tailscale_host
1598     fi
1599
1600     # 2. Wait for physical DHCP lease to settle (crucial for restarts/wake-up/roaming)
1601     wait_for_dhcp_settle
1602
1603     # 3. Boot Colima VM if it is not already running
1604     echo -e "\nChecking Colima VM status..."
1605     if ! run_with_timeout 15 colima status >> output.log 2>&1; then
1606         # Colima networking is gated on the SAME LAN-P2P signal routing-fix uses
1607         # (.lan_p2p_detected overrides TAILSCALE_ALLOW_LAN_P2P in .env). A LAN-reachable
1608         # VM (--network-address + bridged/shared) needs the socket_vmnet helper and is
1609         # only worthwhile on a trusted home network, so we request it only when P2P is
1610         # allowed; otherwise AP-isolated / enterprise, the common case we boot plain
1611         # user-mode networking. Either way host<->container P2P is UNAFFECTED: it rides
1612         # routing-fix's Docker-gateway / .host_ips RETURN rules, not the VM network mode
1613         # (see devref 15.3.5). colima_start_until_ready returns as soon as Docker answers
1614         # (~15s) instead of blocking on colima's ~2min post-provision hang (see 15.8.7);
1615         # that hang is colima-internal and mode-independent (it also occurs in user mode).
1616         allow_p2p=$(read_env_var "TAILSCALE_ALLOW_LAN_P2P")
1617         [ -z "$allow_p2p" ] && allow_p2p="false"
1618         if [ -f "$SCRIPT_DIR/.lan_p2p_detected" ]; then
1619             allow_p2p=$(tr -d '[:space:]' < "$SCRIPT_DIR/.lan_p2p_detected" 2>/dev/null || echo "false")
1620         fi
1621
1622         if [ "$allow_p2p" = "true" ]; then
1623             colima_network_mode="bridged"
1624             if [[ "$OSTYPE" == "darwin*" ]]; then
1625                 security_type=$(system_profiler SPAirPortDataType 2>/dev/null | awk '/Current Network Information
1626                 :/{found=1} found && /Security/{print $NF; exit}')
1627                 if echo "$security_type" | grep -qi "Enterprise"; then
1628                     echo -e "\n [!] WPA2-Enterprise Wi-Fi detected! Bridged mode will cause MAC-spoofing
1629                     deauthentication."
1630                     echo "      Falling back to '--network-mode shared' for Colima."
1631                     colima_network_mode="shared"
1632                 fi

```

```

1630     fi
1631     echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P on: --network-address
$colima_network_mode)... "
1632     colima_start_until_ready --memory 0.6 --vm-type vz --network-address --network-mode "
$colima_network_mode" \
1633     || echo " [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will
retry)."
1634 else
1635     colima_network_mode="user"
1636     echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P off: fast user-mode
networking)... "
1637     colima_start_until_ready --memory 0.6 --vm-type vz \
1638     || echo " [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will
retry)."
1639 fi
1640 echo "$colima_network_mode" > /tmp/nullexit_colima_mode.txt
1641 else
1642     echo "Colima is already running."
1643 fi
1644
1645 # 3b. Auto-detect and fix Docker Subnet Collisions (e.g. if the user travels to a 172.17.x.x Wi-Fi)
1646 if [ -f "scripts/fix-docker-bridge-collision.sh" ]; then
1647     if bash scripts/fix-docker-bridge-collision.sh --dry | grep -q "collision detected"; then
1648         echo -e "\n[!] WARNING: Docker Subnet Collision Detected with your current Wi-Fi!"
1649         echo -e "    Auto-fixing Colima config to prevent internet jitter..."
1650         bash scripts/fix-docker-bridge-collision.sh --skip-toggle || true
1651     fi
1652 fi
1653
1654 # 3c. Configure swap file inside the VM to prevent OOM on low-memory limits
1655 # We also set vm.swappiness=10 so the Linux kernel strictly prefers physical RAM
1656 # and avoids unnecessarily wearing out the SSD with proactive background swapping.
1657 if ! run_with_timeout 15 colima ssh -- grep -q 'swapfile' /proc/swaps >> output.log 2>&1; then
1658     echo "Configuring 400MB swap file inside the VM to prevent OOM..."
1659     run_with_timeout 30 colima ssh -- sudo sh -c "if [ ! -f /swapfile ]; then dd if=/dev/zero of=/
swapfile bs=1M count=400 status=none && mkswap /swapfile; fi && swapon /swapfile && sysctl vm.
swappiness=10" >> output.log 2>&1 || echo "Warning: Failed to enable swap file inside the VM."
1660 fi
1661
1662 # 4. Clean up corrupted AdGuardHome configurations
1663 if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
1664     rm -f "adguard/conf/AdGuardHome.yaml"
1665     echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
1666 fi
1667
1668 # 4b. Auto-heal stale Docker socket from hard reboots
1669 if ! run_with_timeout 15 docker ps >/dev/null 2>&1; then
1670     echo "Docker daemon is unresponsive (likely a stale socket from a crash). Auto-healing Colima..."
1671     run_with_timeout 120 colima restart >> output.log 2>&1
1672 fi
1673
1674 # 4c. Inject TCP MSS clamp from .env into post-rules.txt before starting
1675 if [ -f .env ] && grep -q "^GATEWAY_MSS=" .env; then
1676     GATEWAY_MSS=$(read_env_var GATEWAY_MSS)
1677     if [[ "$GATEWAY_MSS" =~ ^[0-9]+$ ]]; then
1678         # macOS sed syntax
1679         sed -i '' "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null || \
1680         # Linux sed fallback
1681         sed -i "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null
1682     fi
1683 fi
1684
1685 # 5. Start compose services
1686 write_host_ips
1687
1688 # Procedurally generated ports have already been exported at the top of the script.
1689 echo " [Security] Procedurally randomized proxy ports generated to evade fingerprinting."
1690
1691 # --- HONEY-PORT TRIPWIRE ---
1692 # Generate a random trap port. If any malware does a sequential port scan on localhost,
1693 # it will inevitably hit this port. When it does, we instantly fire a macOS alert.
1694 # Reap the honey-port from any previous toggle before arming a new one, so the
1695 # 'nc -l' tripwire never accumulates one idle listener per toggle-on (15.12.20).
1696 pkill -f "nc -l localhost" 2>/dev/null || true
1697 HONEY_PORT=$(get_free_port)
1698 (
1699     if nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
1700         echo "[$(date +%Y-%m-%d %H:%M:%S)] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown
application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$LOG_FILE"
1701         osascript -e 'display notification "An unknown application is aggressively scanning your local
ports. Malware port-scan suspected." with title "nullexit OPSEC Alert" sound name "Basso"' 2>/dev/

```



```

1702     null || true
1703     fi
1704 ) &
1705 disown %1 2>/dev/null || true
1706 echo " [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."
1707
1708 echo -e "\nStarting Docker containers..."
1709
1710 # --- Ad-block rule compilation: hash-gated one-off, off the critical path (15.8.8) ---
1711 # The rule-compiler (a 'compile'-profiled container, the SOLE writer of adguard/work)
1712 # re-dedupes ~340k rules inside the 600MB VM (~28s) far too slow to run every toggle.
1713 # It is gated on a hash of the lists YOU control (black_list.txt + white_list.txt):
1714 # recompile only when they change, when compiled_rules.txt / ip_blocklist.ipset are
1715 # missing, or when they exceed RULES_MAX_AGE_DAYS (default 7, so the remote blocklists
1716 # still refresh occasionally). Otherwise AdGuard/routing-fix boot on the existing
1717 # artifacts and we skip the ~28s they no longer 'depends_on' the rule-compiler
1718 # container (see docker-compose.yml). A compile failure is non-fatal: we keep the
1719 # previous compiled rules rather than bricking the tunnel (ad-block != the kill-switch).
1720 RULES_HASH_FILE="$SCRIPT_DIR/.rules_hash"
1721 rules_hash=$(compute_rules_hash)
1722 stored_rules_hash=""
1723 [ -f "$RULES_HASH_FILE" ] && stored_rules_hash=$(cat "$RULES_HASH_FILE" 2>/dev/null)
1724 rules_max_age_days=$(read_env_var RULES_MAX_AGE_DAYS)
1725 [ -z "$rules_max_age_days" ] && rules_max_age_days=7
1726 need_compile=false
1727 if [ ! -f adguard/work/userfilters/compiled_rules.txt ] || [ ! -f adguard/work/userfilters/ip_blocklist
1728 .ipset ]; then
1729     need_compile=true
1730 elif [ -z "$rules_hash" ] || [ "$rules_hash" != "$stored_rules_hash" ]; then
1731     need_compile=true
1732 elif [ -n "$(find adguard/work/userfilters/compiled_rules.txt -mtime +"$rules_max_age_days" 2>/dev/null
1733 )" ]; then
1734     need_compile=true
1735 fi
1736 if [ "$need_compile" = "true" ]; then
1737     echo " Ad-block lists changed (or artifacts missing/stale) compiling rules (one-off, ~28s)..."
1738     if run_logged docker compose run --build --rm rule-compiler >> output.log 2>&1; then
1739         if [ -n "$rules_hash" ]; then echo "$rules_hash" > "$RULES_HASH_FILE"; fi
1740     else
1741         echo " [!] Rule compile failed continuing with existing compiled ad-block rules."
1742     fi
1743 else
1744     echo " Ad-block lists unchanged reusing compiled rules (skipping the ~28s recompile)."
1745 fi
1746
1747 # --- Gate '--build' for the long-running images (routing-fix, tor) ---
1748 # BuildKit re-checking these on the 600MB VM costs ~30s even when every layer is
1749 # CACHED (15.8.8). Only pass --build when the hashed inputs change or an image is
1750 # missing; otherwise plain 'up -d' reuses cached images. 'up' starts everything
1751 # except the 'compile'-profiled rule-compiler.
1752 BUILD_HASH_FILE="$SCRIPT_DIR/.build_hash"
1753 build_inputs_hash=$(compute_build_inputs_hash)
1754 stored_build_hash=""
1755 [ -f "$BUILD_HASH_FILE" ] && stored_build_hash=$(cat "$BUILD_HASH_FILE" 2>/dev/null)
1756 local_images_present=true
1757 for _img in nullexit-routing-fix:v1.0.0 nullexit-tor:v1.0.0; do
1758     docker image inspect "$_img" >> output.log 2>&1 || local_images_present=false
1759 done
1760 if [ -n "$build_inputs_hash" ] && [ "$build_inputs_hash" = "$stored_build_hash" ] && [ "$local_images_present" = "true" ]; then
1761     echo " Build inputs unchanged skipping image build (reusing cached images)."
1762     run_logged docker compose up -d
1763 else
1764     echo " Build inputs changed or a local image is missing building images..."
1765     run_logged docker compose up -d --build
1766     # Record the hash only after a successful build (set -e aborts above on failure).
1767     if [ -n "$build_inputs_hash" ]; then echo "$build_inputs_hash" > "$BUILD_HASH_FILE"; fi
1768 fi
1769
1770 RULE_COUNT=$(grep "! Total Block Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk
1771 -F': ' '{print $2}' || echo "0")
1772 NATIVE_COUNT=$(grep "! Native AdGuard Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null |
1773 awk -F': ' '{print $2}' || echo "0")
1774
1775 if [ "$RULE_COUNT" != "0" ] && [ "$RULE_COUNT" != "" ]; then
1776     if [ "$NATIVE_COUNT" != "0" ] && [ "$NATIVE_COUNT" != "" ]; then
1777         TOTAL_COUNT=$((RULE_COUNT + NATIVE_COUNT))
1778         # Format with commas for readability (e.g. 441,578)
1779         if command -v printf &> /dev/null; then
1780             FORMATTED_RULE=$(printf "%'d" "$RULE_COUNT")
1781             FORMATTED_TOTAL=$(printf "%'d" "$TOTAL_COUNT")

```

```

1777     echo "   DNS: Compiled $FORMATTED_RULE unique custom rules (Total active in AdGuard: ~
$FORMATTED_TOTAL rules)."
1778     else
1779     echo "   DNS: Compiled $RULE_COUNT unique custom rules (Total active in AdGuard: ~$TOTAL_COUNT
rules)."
1780     fi
1781 else
1782     echo "   DNS: Compiled and loaded $RULE_COUNT optimized rules."
1783 fi
1784 fi
1785
1786 # Report IP blocklist compilation results
1787 IP_COUNT=$(grep "^# Entries:" adguard/work/userfilters/ip_blocklist.ipset 2>/dev/null | awk -F': ' '{
print $2}' || echo "0")
1788 if [ "$IP_COUNT" != "0" ] && [ "$IP_COUNT" != "" ]; then
1789     if command -v printf &> /dev/null; then
1790         FORMATTED_IP=$(printf "%d" "$IP_COUNT")
1791         echo "   IP: Loaded $FORMATTED_IP threat intelligence IPs/CIDRs into kernel firewall."
1792     else
1793         echo "   IP: Loaded $IP_COUNT threat intelligence IPs/CIDRs into kernel firewall."
1794     fi
1795 fi
1796
1797 # 6. Wait for the gateway container's Tailscale connection to be ready
1798 BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
1799 USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')
1800
1801 echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
1802 echo "   (This can take 30-60s Tailscale goes through: NoState Starting Running Online)"
1803 connected=false
1804 consecutive_failures=0
1805 for i in {1..60}; do
1806     # Run tailscale status and capture its output so the user can see what's happening.
1807     # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
1808     ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)
1809
1810     if echo "$ts_output" | grep -q "offers exit node"; then
1811         connected=true
1812         break
1813     fi
1814
1815     # Track consecutive failures to detect a stuck state.
1816     # If we see 40+ consecutive 'NoState' responses, give up early
1817     # the daemon is alive but can't connect to the coordination server.
1818     if echo "$ts_output" | grep -q "NoState"; then
1819         consecutive_failures=$((consecutive_failures + 1))
1820         if [ "$consecutive_failures" -ge 40 ]; then
1821             echo ""
1822             echo "   [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
1823             echo "   Tailscale daemon is running but can't reach the coordination server."
1824             break
1825         fi
1826     else
1827         consecutive_failures=0
1828     fi
1829
1830     # Print a condensed status every 10 seconds so the user can see progress
1831     if [ $((i % 10)) -eq 0 ]; then
1832         ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
1833         echo "   [attempt $i/60] $ts_short"
1834     elif [ $((i % 5)) -eq 0 ]; then
1835         echo -n "$i"
1836     else
1837         echo -n " ."
1838     fi
1839     sleep 1
1840 done
1841 echo ""
1842
1843 if [ "$connected" = "true" ]; then
1844     echo "Gateway container is online on the Tailscale mesh."
1845 elif [ "$BYPASS_PING" = "true" ]; then
1846     echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
1847 else
1848     echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
1849     echo "   The container was seen in the following states during the wait:" >&2
1850     echo "$ts_output" | sed 's/^/    /' >&2
1851     echo "" >&2
1852     echo "   This could mean:" >&2
1853     echo "       1. Tailscale auth key has expired check TS_AUTHKEY in .env" >&2

```

```

1854 echo "      2. Network issue inside the container  check: docker compose logs tailscale" >&2
1855 echo "      3. gluetun VPN tunnel not connecting  check: docker compose logs warp | tail -20" >&2
1856 echo "" >&2
1857 echo "      Diagnose manually:" >&2
1858 echo "      docker compose exec -T tailscale tailscale status" >&2
1859 echo "      docker compose exec -T tailscale tailscale netcheck" >&2
1860 cleanup_handler ERR $LINENO
1861 exit 1
1862 fi
1863
1864 # 6b. Resolve the gateway's Tailscale IP.
1865 # Dynamically query the container for the IP (handles IP changes after re-auth)
1866 # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
1867 #
1868 # CRITICAL: 'docker compose exec -T' on macOS/Colima injects carriage
1869 # returns (\r) into captured output. If $TS_IP ends with \r, then
1870 # 'networksetup -setdnsservers' silently rejects the malformed IP and
1871 # DNS stays at 1.1.1.1 the primary bug this project was hitting.
1872 # 'tr -d '\r'' strips the CR; 'awk 'NR==1{print $1; exit}'' takes only
1873 # the first field of the first line.
1874 TS_IP=""
1875 if [ "$connected" = "true" ]; then
1876     TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
1877     if [ -n "$TS_IP" ]; then
1878         echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
1879         echo "$TS_IP" > .gateway_ip
1880     else
1881         echo "      [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
1882         TS_IP=$(read_adguard_ip || true)
1883     fi
1884 else
1885     TS_IP=$(read_adguard_ip || true)
1886 fi
1887
1888 # 7. Connect host Mac to Tailscale mesh.
1889 # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
1890 # system LaunchDaemon via 'brew services start tailscale'), the previous five-
1891 # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
1892 # no menu bar item to click, and no Accessibility permission required.
1893 #
1894 # CRITICAL: If tailscaled is not running, 'tailscale up' will silently fail.
1895 # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
1896 # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
1897 # are impossible without the host on the mesh).
1898 #
1899 # We connect in two phases:
1900 #   Phase A Join mesh WITHOUT exit node (so pre-flight checks can run)
1901 #   Phase B Set exit node only after pre-flight checks confirm the gateway works
1902 HOST_ON_MESH=false
1903 SKIP_EXIT_NODE=false
1904 if [ "$HIJACK_HOST" = "false" ]; then
1905     echo -e "\n[Info] HIJACK_HOST is false (Headless Mode). Host networking will remain untouched."
1906     SKIP_EXIT_NODE=true
1907 elif [ -n "$TS_BIN" ]; then
1908     echo -n "Verifying tailscaled is reachable"
1909     daemon_ready=false
1910     status_exit=0
1911     for i in {1..10}; do
1912         # Test if the daemon responds to status check.
1913         # If status returns non-zero, check if it was a normal "stopped/disconnected" state
1914         # (exit code 1 is normal for disconnected). If the command exited quickly (not killed by timeout),
1915         # and the error is just "Tailscale is stopped", then the daemon is alive and ready.
1916         # But if it timed out (exit code 143) or is completely unresponsive, it's wedged.
1917         status_exit=0
1918         if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
1919             daemon_ready=true
1920             break
1921         else
1922             status_exit=$?
1923             if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
1924                 daemon_ready=true
1925                 break
1926             fi
1927         fi
1928         echo -n "."
1929         sleep 1
1930     done
1931     echo ""
1932
1933     # If the daemon was unresponsive or socket check failed, attempt auto-restart

```

```

1934 if [ "$daemon_ready" != "true" ]; then
1935     if [ "$status_exit" -eq 143 ]; then
1936         warn "tailscaled daemon is wedged/unresponsive. Restarting it..."
1937     else
1938         warn "tailscaled daemon is not running. Attempting to start it..."
1939     fi
1940     restart_tailscaled_daemon
1941
1942     # Re-verify
1943     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
1944         daemon_ready=true
1945     fi
1946 fi
1947
1948 if [ "$daemon_ready" = "true" ]; then
1949     echo "tailscaled is responsive (running as a system service)."
1950
1951     # Phase A: Join mesh without exit node
1952     echo "Connecting host to Tailscale mesh (no exit node yet)..."
1953     if $TS_BIN up; then
1954         $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node= || true
1955         HOST_ON_MESH=true
1956         echo "Host is on Tailscale mesh."
1957     else
1958         warn "tailscale up failed even though the daemon is running."
1959         echo "  This usually means the host hasn't authenticated with your tailnet yet."
1960         echo "  Run this command in a terminal to authenticate:"
1961         echo ""
1962         echo "      sudo tailscale up"
1963         echo ""
1964         echo "  A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
1965     fi
1966 else
1967     fail "tailscaled could NOT be started automatically."
1968     echo "  Run this in a terminal to start and authenticate Tailscale:"
1969     echo ""
1970     echo "      sudo brew services start tailscale"
1971     echo "      sudo tailscale up"
1972     echo ""
1973     echo "  Then re-run toggle.sh."
1974 fi
1975 else
1976     echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
1977 fi
1978
1979 # Phase B: Pre-flight checks + enable exit node only if safe
1980 if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
1981     echo -e "\n"
1982     echo "Pre-flight connectivity check"
1983     echo ""
1984
1985     # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
1986     # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
1987     # ("direct connection not established") even though the pong was received.
1988     # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
1989     echo -n " [1/3] Gateway reachable via Tailscale... "
1990     ping_ok=false
1991     # The host tailscaled needs a few seconds to propagate the new network map and
1992     # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
1993     for attempt in {1..45}; do
1994         if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
1995             ping_ok=true
1996             break
1997         fi
1998         sleep 1
1999     done
2000     if [ "$ping_ok" = "true" ]; then
2001         echo "PASS"
2002     else
2003         echo "FAIL"
2004         SKIP_EXIT_NODE=true
2005     fi
2006
2007     # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
2008     # Colima's SSH tunnel only forwards TCP (not UDP), so we use +tcp.
2009     echo -n " [2/3] AdGuard DNS via localhost... "
2010     if command -v dig &>/dev/null; then
2011         if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 &>/dev/null; then
2012             echo "PASS"
2013         else
2014             echo "FAIL"

```

```

2015     SKIP_EXIT_NODE=true
2016 fi
2017 elif command -v nslookup &>/dev/null; then
2018     if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 &>/dev/null; then
2019         echo "PASS"
2020     else
2021         echo "FAIL"
2022         SKIP_EXIT_NODE=true
2023     fi
2024 else
2025     echo "SKIP (dig/nslookup not available)"
2026 fi
2027
2028 # Check 3: Does the WARP container have external internet?
2029 echo -n " [3/3] WARP container internet... "
2030 tmp_err=$(mktemp)
2031 if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace >/dev/
null 2>"$tmp_err"; then
2032     echo "PASS"
2033 else
2034     echo "FAIL"
2035     if [ -s "$tmp_err" ]; then
2036         err_msg=$(cat "$tmp_err" | tr -d '\r')
2037         echo "[$(date +%Y-%m-%d %H:%M:%S)] [Check 3] WARP container check failed: $err_msg" >> output.
log
2038     fi
2039     SKIP_EXIT_NODE=true
2040 fi
2041 rm -f "$tmp_err"
2042
2043 # Act based on check results
2044 if [ "$SKIP_EXIT_NODE" != "true" ]; then
2045     echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
2046     $TS_BIN up || true
2047     if $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node="$TS_IP" --exit-node-
allow-lan-access=true; then
2048         echo "Exit node enabled."
2049         add_warp_bypass_routes
2050         setup_exit_node_routing
2051         enable_killswitch
2052         # Opt-in FaceTime split-route (no-op unless FACETIME_SPLIT_ROUTE=true).
2053         # After the kill-switch so the com.apple/nullexit anchor + table exist.
2054         add_apple_split_routes
2055     else
2056         echo "[Warning] Failed to set exit node."
2057         SKIP_EXIT_NODE=true
2058     fi
2059 fi
2060
2061 if [ "$SKIP_EXIT_NODE" = "true" ]; then
2062     warn "Pre-flight checks failed EXIT NODE + DNS HIJACK SKIPPED."
2063     echo " Host stays on Tailscale mesh but your internet routes through your normal connection."
2064     echo " Troubleshoot with:"
2065     echo "     docker compose logs warp | tail -20"
2066     echo "     docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/
trace"
2067     echo ""
2068     echo " Falling back to SOCKS5 proxy for traffic routing..."
2069 fi
2070 elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
2071     echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
2072     SKIP_EXIT_NODE=true
2073 fi
2074
2075 # 8. Apply DNS Hijacking.
2076 # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
2077 # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
2078 # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
2079 if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
2080     echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
2081     echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
2082     if force_dns_to_gateway "$TS_IP"; then
2083         echo "DNS hijack verified: single server = $TS_IP."
2084     else
2085         echo "[Warning] DNS hijack could not be verified."
2086         echo " If DNS is broken, run manually:"
2087         echo "     networksetup -setdnsservers \"\$ACTIVE_SERVICE\" \"$TS_IP"
2088         echo "     networksetup -setsearchdomains \"\$ACTIVE_SERVICE\" \"ts.net"
2089     fi
2090 elif [ -n "$TS_IP" ] && [ "$HIJACK_HOST" != "false" ]; then
2091     echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."

```

```

2092 echo " Trying local DNS proxy for ad-blocking..."
2093 if start_local_dns_proxy; then
2094     echo " Ad-blocking active via local DNS proxy through AdGuard."
2095 else
2096     echo " Local DNS proxy unavailable. DNS stays at 1.1.1.1."
2097     echo " AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
2098 fi
2099
2100 # Enable SOCKS5 proxy for traffic routing through WARP
2101 # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
2102 echo -e "\n Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
2103 echo " (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
2104 enable_socks_proxy "$ACTIVE_SERVICE"
2105
2106 # Verify the SOCKS5 proxy works through WARP
2107 echo -n " Verifying SOCKS5 proxy routes through WARP... "
2108 if curl --socks5-hostname 127.0.0.1:${SOCKS_PROXY_PORT} --max-time 10 -s https://www.cloudflare.com/cdn
-cgi/trace 2>/dev/null | grep -q "warp=on"; then
2109     echo "ok (warp=on)."
2110     echo " All TCP traffic now encrypted through Cloudflare WARP tunnel."
2111 else
2112     echo "FAILED (proxy not routing through WARP)."
2113     echo " Disabling SOCKS5 proxy internet stays on direct connection."
2114     disable_socks_proxy
2115     echo " SOCKS proxy available at localhost:${SOCKS_PROXY_PORT} for manual configuration."
2116 fi
2117 else
2118     echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
2119 fi
2120
2121 # 9. Verify host connectivity
2122 if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
2123     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
2124         echo "Host is online on the Tailscale mesh."
2125     else
2126         echo "[Warning] Host is not responding on Tailscale."
2127     fi
2128 fi
2129
2130 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
2131 echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."
2132
2133 # Prevent system from going to sleep while gateway is running
2134 start_sleep_prevention
2135
2136 if [ "$HIJACK_HOST" != "false" ]; then
2137     if [ -n "$TS_IP" ] && [ "$TS_IP" != "1.1.1.1" ]; then
2138         start_dns_watcher "$TS_IP"
2139     fi
2140
2141     # Start WARP liveness monitor (logs to output.log on warp flip)
2142     start_warp_watcher
2143 fi
2144
2145 # Reset sharing services after network configuration to prevent AirDrop freezes
2146 echo "Resetting macOS sharing services (AirDrop/AirPlay)..."
2147 reset_sharing_services
2148
2149 # Tell external watchers (post-wake / network-change) the gateway is up.
2150 write_gateway_active_marker
2151
2152 # Local Network Surveillance Check
2153 echo -e "\n"
2154 echo "Local Network Surveillance Check"
2155 echo ""
2156 # Count populated ARP cache entries to detect network scanning or high congestion
2157 # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
2158 ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
2159 if [ "$ARP_COUNT" -gt 15 ]; then
2160     warn "High ARP activity detected ($ARP_COUNT devices in cache)."
2161     warn "This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)."
2162     echo -e " [] Your traffic remains fully encrypted and invisible.\n"
2163 else
2164     echo -e " [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
2165 fi
2166 fi
2167
2168 SUCCESS_RUN=true
2169
2170 # Final DNS state summary
2171 if [ -n "$ACTIVE_SERVICE" ]; then

```



```

2172 FINAL_DNS=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log || true)
2173 echo -e "\n"
2174 echo -e "DNS STATE: $FINAL_DNS"
2175 echo -e ""
2176 fi
2177
2178 if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
2179     if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
2180         echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
2181         echo " The gateway gracefully fell back to yesterday's cached blocklist."
2182         echo " (Run 'docker logs rule-compiler' later to investigate which link died.)"
2183     fi
2184 fi
2185
2186 rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
2187 if [ -t 0 ]; then
2188     read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
2189     if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
2190         echo "Reversing gateway state..."
2191         exec bash "$0"
2192     fi
2193 fi
2194
2195 echo -e "\nYou can close this terminal window now."

```

7 recover.sh

```

1 #!/bin/bash
2 # recover.sh nullexit KILL-SWITCH recovery script (macOS + Linux)
3 # Fixes internet when toggle.sh leaves you stranded.
4 #
5 # This is a NUCLEAR option: it undoes EVERYTHING toggle.sh does by:
6 # 1. Disconnecting host Tailscale from the mesh (exit-node off)
7 # 2. Resetting DNS to default on ALL network services
8 # 3. Disabling rogue proxy settings (SOCKS, web, secure web)
9 # 4. Flushing the DNS cache
10 # 5. Flushing stale routing table entries
11 # 6. Stopping gateway Docker containers
12 # 7. Power-cycling Wi-Fi
13 # 8. Verifying internet is working
14 #
15 # One file, both OSes: the dispatch below runs the self-contained Linux
16 # recovery and exits; on macOS it falls through to the richer dual-mode
17 # implementation (default teardown / --post-wake refresh).
18 #
19 # Usage: double-click Recover-Gateway.applescript
20 #         OR ./recover.sh
21 #         OR ./recover.sh --post-wake (macOS only lighter refresh, keeps gateway live)
22
23 set -e
24
25 # Resolve script directory (used for path-relative refs)
26 # recover.sh lives at the repo root; SCRIPT_DIR lets us launch toggle.sh
27 # (and reference any other repo-root file) from the same directory no
28 # matter where the user invoked recover.sh from.
29 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
30 cd "$SCRIPT_DIR" || exit 1
31
32 #
33 # LINUX RECOVERY self-contained; runs the full teardown then exits before the
34 # macOS body below. Native Linux tooling (resolvectl / nmcli / ip) throughout.
35 #
36 if [[ "$OSTYPE" != "darwin"* ]]; then
37     # Source common formatting and helper functions
38     source "$SCRIPT_DIR/scripts/common.sh"
39
40     # Always add Homebrew path
41     export PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:$PATH"
42
43     echo ""
44     echo -e "${RED}${BOLD}${NC}"
45     echo -e "${RED}${BOLD} Nullexit Gateway RECOVERY MODE ${NC}"
46     echo -e "${RED}${BOLD}${NC}"
47     echo ""
48     echo "This script will tear down the gateway and restore normal internet."
49     echo ""

```

```

50
51 # Cache sudo credentials upfront
52 # Prevents the script from hanging halfway through when it needs to run
53 # privileged commands (route flush, Wi-Fi power cycle, etc.).
54 echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
55 sudo -v
56 (
57     while true; do
58         sudo -n true 2>/dev/null
59         sleep 60
60     done
61 ) &
62 SUDO_KEEPER_PID=$!
63 trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
64
65 # 1. Disconnect host Tailscale
66 step "Disconnecting host Tailscale from mesh"
67 if command -v tailscale >> output.log 2>&1; then
68     # Check if actually connected first (with timeout tailscaled could be wedged)
69     if run_with_timeout 5 tailscale status >> output.log 2>&1; then
70         # Also explicitly reset any exit-node so it doesn't linger.
71         if run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= >> output.
72             log 2>&1; then
73             ok "Exit-node preference cleared"
74         else
75             warn "tailscale up --reset didn't respond (tailscaled may be wedged)"
76         fi
77         ts_args=""
78         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
79         if run_with_timeout 10 tailscale down $ts_args >> output.log 2>&1; then
80             ok "Tailscale disconnected"
81         else
82             warn "tailscale down didn't respond"
83         fi
84     else
85         warn "tailscaled not reachable skipping (timeout after 5s)"
86     fi
87 else
88     warn "tailscale CLI not found skipping"
89 fi
90
91 # 2. Reset DNS to default on ALL network services
92 step "Resetting DNS to default on active interface"
93
94 ACTIVE_SERVICE=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
95 if [ -n "$ACTIVE_SERVICE" ]; then
96     sudo resolvectl revert "$ACTIVE_SERVICE" 2>> output.log || true
97     ok "DNS reset on $ACTIVE_SERVICE"
98 else
99     warn "Could not determine active network interface"
100 fi
101
102 # 3. Disable rogue proxy settings
103 step "Disabling proxy settings (SOCKS, web, secure web)"
104 echo " (Linux global SOCKS proxy configuration skipped.)"
105 ok "Proxy settings cleared"
106
107 # 4. Flush DNS cache
108 step "Flushing DNS cache"
109 if command -v resolvectl >> output.log 2>&1; then
110     sudo resolvectl flush-caches 2>> output.log || true
111     ok "DNS cache flushed"
112 else
113     warn "resolvectl not found"
114 fi
115
116 # 5. Clear stale routing table
117 step "Flushing stale routing table entries"
118 sudo ip route flush cache >> output.log 2>&1 || true
119 ok "Routing table flushed"
120
121 # 6. Stop gateway Docker containers
122 step "Stopping gateway Docker containers"
123 if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
124     if [ -f "docker-compose.yml" ]; then
125         docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were
126             running"
127     else
128         warn "docker-compose.yml not found in current directory"
129     fi

```

```

129 else
130     warn "Docker not available skipping"
131 fi
132
133 # 6b. Stopping sleep prevention (systemd-inhibit)
134 step "Stopping sleep prevention"
135 PID_FILE="/tmp/nullexit-caffeinate.pid"
136 if [ -f "$PID_FILE" ]; then
137     INHIBIT_PID=$(cat "$PID_FILE")
138     if [ -n "$INHIBIT_PID" ] && kill -0 "$INHIBIT_PID" 2>/dev/null; then
139         kill "$INHIBIT_PID" 2>/dev/null || true
140         ok "Sleep prevention stopped (PID $INHIBIT_PID)"
141     else
142         warn "Sleep prevention process not running, cleaning up stale PID file"
143     fi
144     rm -f "$PID_FILE"
145 else
146     ok "Sleep prevention was not active"
147 fi
148
149 # 7. Power-cycle Wi-Fi
150 step "Power-cycling Wi-Fi"
151 if command -v nmcli &>/dev/null; then
152     sudo nmcli radio wifi off 2>> output.log || true
153     sleep 2
154     sudo nmcli radio wifi on 2>> output.log || true
155     ok "Wi-Fi bounced via nmcli"
156     sleep 3
157 else
158     warn "nmcli not found. Falling back to ip link..."
159     WIFI_IFACE=$(ip link | grep -E '^[0-9]+: (wl[a-zA-Z0-9]+|wlan[0-9]+):' | awk -F': ' '{print $2}')
160     if [ -n "$WIFI_IFACE" ]; then
161         sudo ip link set "$WIFI_IFACE" down 2>> output.log || true
162         sleep 2
163         sudo ip link set "$WIFI_IFACE" up 2>> output.log || true
164         ok "Wi-Fi bounced via ip link (interface $WIFI_IFACE)"
165         sleep 3
166     else
167         warn "No Wi-Fi interface detected."
168     fi
169 fi
170
171 # 8. Verify internet connectivity
172 verify_internet_connectivity
173
174 # Summary
175 echo ""
176 echo -e "${BOLD}${NC}"
177 echo -e "${BOLD}RECOVERY SUMMARY${NC}"
178 echo -e "${BOLD}${NC}"
179
180 # Show final DNS state
181 FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+'
182 | tr '\n' ' ' || echo "N/A")
183 echo " DNS ($ACTIVE_SERVICE): $FINAL_DNS"
184
185 echo ""
186
187 if [ "$INTERNET_OK" = "true" ]; then
188     echo -e " ${GREEN}${BOLD} Internet is working.${NC}"
189 else
190     echo -e " ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
191     echo " This might mean:"
192     echo " - Your Wi-Fi is disconnected from the router"
193     echo " - Tailscale is still interfering (try restarting tailscaled:"
194     echo "     systemctl restart tailscaled)"
195     echo " - You need to re-connect to Wi-Fi in the menu bar"
196     echo " - Try toggling Wi-Fi off/on in the menu bar"
197     echo ""
198     echo " Or try manually:"
199     echo "     resolvectl revert $ACTIVE_SERVICE"
200     extra_args=""
201     if is_kill_switch_enabled; then extra_args="--accept-risk=lose-ssh"; fi
202     echo "     tailscale down $extra_args"
203     echo "     sudo ip route flush cache"
204     echo "     systemctl restart tailscaled"
205 fi
206
207 echo ""
208 echo -e "${GREEN}${BOLD} Recovery complete.${NC}"
209 echo ""

```

```

209     exit 0
210 fi
211 #
212 #
213 # macOS RECOVERY (dual-mode: default teardown / --post-wake refresh)
214 #
215 #
216 # Enforce Cryptographic Script Integrity
217 if [ -f "scripts/crypto.sh" ]; then
218     if ! bash scripts/crypto.sh --verify; then
219         exit 1
220     fi
221 fi
222 #
223 #
224 # Argument parsing
225 # --post-wake: lightweight refresh after sleep/wake or Wi-Fi roam.
226 #
227 #     Keeps the gateway live while HUPing DNS caches, refreshing the
228 #     Tailscale exit-node leak, re-hijacking host DNS, and force-
229 #     recreating the warp container if its gluetun healthcheck failed.
230 #
231 #     Does NOT tear down Docker, Tailscale, sleep prevention, or Wi-Fi.
232 #     Called from scripts/watcher.sh on every wake / network-change event
233 #     while /tmp/nullexit-gateway-active.marker is present (i.e. the
234 #     gateway is currently up). See devref.md 10.29 for the why.
235 POST_WAKE=false
236 AUTO_MODE=false
237 for arg in "$@"; do
238     if [ "$arg" = "--post-wake" ]; then
239         POST_WAKE=true
240     elif [ "$arg" = "--auto" ] || [ "$arg" = "--non-interactive" ]; then
241         AUTO_MODE=true
242     fi
243 done
244 rm -f "$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker"
245 # Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
246 LOCK_FILE="/tmp/nullexit-toggle.lock"
247 if [ -f "$LOCK_FILE" ]; then
248     LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
249     if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
250         if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
251             if [ "$POST_WAKE" = "true" ]; then
252                 echo "[(date -u +%FT%TZ)] POST-WAKE skipped: another lifecycle script (PID $LOCK_PID) is running"
253                 . " >> output.log"
254                 exit 0
255             else
256                 echo "Error: Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running." >&2
257                 exit 1
258             fi
259         fi
260     fi
261     echo "$$" > "$LOCK_FILE"
262 # Clear the active-state marker in nuclear recovery mode since the gateway is being
263 # completely torn down. This prevents scripts/watcher.sh from triggering post-wake
264 # recoveries in response to network changes caused by the teardown itself.
265 if [ "$POST_WAKE" = "false" ]; then
266     rm -f "/tmp/nullexit-gateway-active.marker"
267 fi
268 #
269 #
270 # WARP Watcher inhibit marker: written immediately in --post-wake mode so the
271 # background WARP Watcher (in toggle.sh) skips failure counting while we are
272 # intentionally tearing down / recreating the warp container. Without this, the
273 # watcher accumulates 6 consecutive warp=off readings during force-recreate (~35s)
274 # and fires nuclear recover.sh, killing a gateway that would have been healthy.
275 WARP_INHIBIT_FILE="/tmp/nullexit-warp-inhibit.marker"
276 if [ "$POST_WAKE" = "true" ]; then
277     touch "$WARP_INHIBIT_FILE"
278     echo "[(date -u +%FT%TZ)] POST-WAKE started WARP Watcher inhibited to prevent spurious shutdown
279         during warp recreate." >> output.log
280 fi
281 # Ensure the inhibit marker and lock file are always removed when the script exits (success or error)
282 cleanup_recover() {
283     rm -f "$WARP_INHIBIT_FILE"
284     if [ -f "$LOCK_FILE" ] && [ "$(cat "$LOCK_FILE" 2>/dev/null)" = "$$" ]; then
285         rm -f "$LOCK_FILE"
286     fi
287 }

```

```

288 trap cleanup_recover EXIT
289
290 # Source common formatting and helper functions
291 source "$SCRIPT_DIR/scripts/common.sh"
292
293 # Always add Homebrew path
294 setup_standard_path
295
296 echo ""
297 if [ "$POST_WAKE" = "true" ]; then
298     echo -e "${YELLOW}${BOLD}${NC}"
299     echo -e "${YELLOW}${BOLD} Nullexit POST-WAKE / POST-ROAM LIGHT ${NC}"
300     echo -e "${YELLOW}${BOLD} (keeps the gateway live, refreshes) ${NC}"
301     echo -e "${YELLOW}${BOLD}${NC}"
302     echo ""
303     echo "Re-hijacking DNS, refreshing Tailscale exit-node, force-recreating"
304     echo "warp if unhealthy, and resetting sharing services. The gateway stays"
305     echo "up. docker compose / Wi-Fi / caffeinate are NOT touched."
306     echo ""
307 else
308     echo -e "${RED}${BOLD}${NC}"
309     echo -e "${RED}${BOLD} Nullexit Gateway RECOVERY MODE ${NC}"
310     echo -e "${RED}${BOLD}${NC}"
311     echo ""
312     echo "This script will tear down the gateway and restore normal internet."
313     echo ""
314 fi
315
316 # Cache sudo credentials upfront (default mode ONLY)
317 # Skip in --post-wake: the LaunchAgent-launched watcher has no TTY, so 'sudo -v'
318 # would block forever waiting for a password. All post-wake commands use 'sudo -n'
319 # which is non-blocking and relies on /etc/sudoers.d/nullexit NOPASSWD entries.
320 SUDO_KEEPER_PID=""
321 if [ "$POST_WAKE" = "false" ] && [ "$AUTO_MODE" = "false" ] && [ -t 0 ]; then
322     echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
323     sudo -v
324     (
325         while true; do
326             sudo -n true 2>/dev/null
327             sleep 60
328         done
329     ) &
330     SUDO_KEEPER_PID=$!
331     trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
332 fi
333
334 # Resolve physical default interface and gateway upfront
335 PHYSICAL_IFACE=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)')
336 [ -z "$PHYSICAL_IFACE" ] && PHYSICAL_IFACE="en0"
337
338 if [ "$POST_WAKE" = "true" ]; then
339     wait_for_dhcp_settle
340 else
341     # Quick resolution in non-post-wake mode (non-blocking)
342     PHYSICAL_GW=$(ipconfig getpacket "$PHYSICAL_IFACE" 2>> output.log | awk -F'[{}]' '/router/{print $2}')
343 fi
344
345 # 1. Disconnect host Tailscale (default) / Refresh exit-node (post-wake)
346
347 if [ "$POST_WAKE" = "true" ]; then
348     step "Refreshing host Tailscale exit-node routing (post-wake)"
349     if command -v tailscale >> output.log 2>&1; then
350         # Re-issue the SAME exit-node up command toggle.sh START uses. This refreshes
351         # the DERP relay mapping and re-asserts the exit-node preference without
352         # dropping the host's mesh connection (which 'tailscale down' would).
353         TS_IP=""
354         for src in .gateway_ip; do
355             [ -f "$src" ] || continue
356             TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
357             [ -n "$TS_IP" ] && break
358         done
359
360         if [ -n "$TS_IP" ]; then
361             step "Re-establishing exit node $TS_IP via 'tailscale set --exit-node'"
362             # 'tailscale set' only mutates the named preference (exit-node here).
363             # Crucially it does NOT call --reset, which would re-apply ALL flags and
364             # trigger a sub-second route-table re-evaluation cycle each post-wake.
365             # On already-connected nodes, the lighter 'set' keeps the existing tunnel
366             # alive while only changing the exit-node preference.
367             if run_with_timeout 15 tailscale set --exit-node="$TS_IP" \

```

```

368             --exit-node-allow-lan-access=true \
369             >> output.log 2>&1; then
370         ok "Exit-node $TS_IP re-asserted via 'tailscale set'"
371     else
372         warn "tailscale set --exit-node=$TS_IP didn't respond; falling back to mesh-only"
373         run_with_timeout 10 tailscale set --exit-node= \
374             >> output.log 2>&1 || true
375     fi
376 else
377     warn ".gateway_ip missing cannot re-assert exit node"
378     run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= \
379         >> output.log 2>&1 || true
380 fi
381 else
382     warn "tailscale CLI not found"
383 fi
384 else
385     step "Disconnecting host Tailscale from mesh"
386     disconnect_tailscale_host "tailscale"
387 fi
388
389 # 2. Reset DNS to default on ALL network services (default) / Re-hijack gateway DNS (post-wake)
390 if [ "$POST_WAKE" = "true" ]; then
391     step "Re-hijacking host DNS to gateway IP (post-wake / post-roam)"
392     TS_IP=""
393     for src in .gateway_ip; do
394         [ -f "$src" ] || continue
395         TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
396         [ -n "$TS_IP" ] && break
397     done
398
399     if [ -n "$TS_IP" ]; then
400         # Detect the active network service (using helper in common.sh)
401         ACTIVE_SVC=$(get_active_service)
402         ENO_SVC=$(get_en0_service)
403
404         networksetup -setsearchdomains "$ACTIVE_SVC" "ts.net" 2>> output.log || true
405         networksetup -setdnsservers "$ACTIVE_SVC" "$TS_IP" 2>> output.log || true
406         networksetup -setsearchdomains "$ENO_SVC" "ts.net" 2>> output.log || true
407         networksetup -setdnsservers "$ENO_SVC" "$TS_IP" 2>> output.log || true
408
409         HITS=$(networksetup -getdnsservers "$ACTIVE_SVC" 2>> output.log | grep -Fx "$TS_IP" || true)
410         if [ -n "$HITS" ]; then
411             ok "DNS re-hijacked to $TS_IP on services: $ACTIVE_SVC, $ENO_SVC"
412         else
413             warn "networksetup accepted but DNS read-back didn't include $TS_IP"
414         fi
415     else
416         warn ".gateway_ip missing cannot re-hijack DNS (user must run toggle.sh START)"
417     fi
418 else
419     step "Resetting DNS to default on all network services (empty = DHCP)"
420
421     # Get ALL network services (handles spaces in names, skips disabled ones)
422     SERVICES=$(networksetup -listallnetworkservices 2>> output.log | grep -v "An asterisk" | grep -v "^$"
423         || true)
424
425     if [ -z "$SERVICES" ]; then
426         # Fallback to common names
427         SERVICES="Wi-Fi Ethernet Thunderbolt Ethernet USB 10/100 LAN"
428     fi
429
430     RESET_COUNT=0
431     while IFS= read -r service; do
432         # Strip leading asterisk (disabled services)
433         service_clean=$(echo "$service" | sed 's/^\*//')
434         [ -z "$service_clean" ] && continue
435
436         # Check if this service has a hardware port (skip VPN/service-only entries)
437         if networksetup -listallhardwareports 2>> output.log | grep -A1 "Port: $service_clean$ | grep -q "
438             Device:" || [ "$service_clean" = "Wi-Fi" ]; then
439             (
440                 networksetup -setsearchdomains "$service_clean" "Empty" 2>> output.log
441                 # Set to "empty" so macOS uses DHCP-assigned DNS (not hardcoded 1.1.1.1)
442                 sudo networksetup -setdnsservers "$service_clean" "empty" 2>> output.log
443             ) && RESET_COUNT=$((RESET_COUNT + 1)) || true
444         fi
445     done <<< "$SERVICES"
446
447     ok "DNS reset on $RESET_COUNT service(s)"
448 fi

```



```

447
448 # 3. Disable rogue proxy settings (default only gateway uses proxy in non-exit-node path)
449 if [ "$POST_WAKE" = "false" ]; then
450     step "Disabling proxy settings (SOCKS, web, secure web)"
451     for svc in $(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "$
        " || echo "Wi-Fi"); do
452         svc_clean=$(echo "$svc" | sed 's/~\*///')
453         [ -z "$svc_clean" ] && continue
454         disable_all_proxies "$svc_clean"
455     done
456     ok "Proxy settings cleared"
457 else
458     step "Proxy settings: leaving untouched (post-wake keeps gateway SOCKS wiring alive)"
459 fi
460
461 # 4. Flush macOS DNS cache (always safe and useful in both modes)
462 step "Flushing DNS cache"
463 if command -v dscacheutil >> output.log 2>&1; then
464     flush_dns_cache
465     ok "DNS cache flushed (mDNSResponder HUP)"
466 else
467     warn "dscacheutil not found"
468 fi
469
470 # 5. Clear stale routing table (always safe in both modes)
471 step "Flushing stale routing table entries"
472 # Resolve physical default gateway IP before flushing
473 # We cannot trust 'route get default' here because Tailscale may have already hijacked
474 # the default route to utunX. We must read the actual DHCP router assignment from the hardware.
475 # PHYSICAL_IFACE and PHYSICAL_GW resolved upfront
476
477 # Instead of full route -n flush (which destroys macOS loopback/multicast and wedges the OS until reboot)
478 # we cleanly delete the Tailscale overrides and Cloudflare WARP bypass routes in ALL modes.
479 # This ensures tailscaled isn't blocked from reaching its control plane when waking from sleep.
480 sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
481 sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
482 remove_warp_bypass_routes
483 ok "Routing overrides cleared"
484
485 if [ "$POST_WAKE" = "false" ]; then
486     disable_killswitch
487     ok "Routing table flush completed via targeted deletes and firewall disabled"
488 else
489     ok "Routing overrides cleared (post-wake mode to preserve Colima bridge100)"
490 fi
491
492 # In post-wake mode we don't bounce Wi-Fi, so we MUST manually restore the physical default route
493 if [ "$POST_WAKE" = "true" ] && [ -n "$PHYSICAL_GW" ]; then
494     # Ensure physical gateway is set just in case it was dropped
495     sudo -n route add default "$PHYSICAL_GW" >> output.log 2>&1 || true
496 fi
497
498 # CONDITIONAL BYPASS ROUTE RESTORATION:
499 # If this was a post-wake recovery and the exit node is active, the route flush
500 # just wiped the static bypass routes for the WARP endpoints. We must re-add
501 # them immediately to prevent a routing loop when Tailscale starts sending traffic.
502 if [ "$POST_WAKE" = "true" ]; then
503     if [ -z "${ACTIVE_IF:-}" ]; then
504         ACTIVE_IF=$(route get default 2>> output.log | awk '/interface:/{print $2; exit}')
505         if [ -z "$ACTIVE_IF" ] || [[ "$ACTIVE_IF" =~ ^utun ]] || [[ "$ACTIVE_IF" =~ ^tun ]]; then
506             for i in en0 en1 en2 en3; do
507                 if ifconfig "$i" 2>> output.log | grep -q 'status: active'; then
508                     ACTIVE_IF="$i"; break
509                 fi
510             done
511         fi
512         [ -z "$ACTIVE_IF" ] && ACTIVE_IF="en0"
513     fi
514
515     echo "Re-adding host bypass routes for Cloudflare WARP endpoints..."
516     remove_warp_bypass_routes
517     if [ -n "${PHYSICAL_GW:-}" ]; then
518         add_warp_bypass_routes "$PHYSICAL_GW"
519     else
520         add_warp_bypass_routes "$ACTIVE_IF"
521     fi
522
523     # Re-assert the opt-in FaceTime split-route so it survives roam/wake
524     # (no-op unless FACETIME_SPLIT_ROUTE=true). Physical gateway is known here.
525     add_apple_split_routes "${PHYSICAL_GW:-}"

```

```

526 # Also re-apply the default route pointing to the Tailscale interface
527 ts_iface=$(ifconfig 2>> output.log | grep -B4 "inet 100." | grep -E '[a-z0-9]+' | cut -d: -f1 | head -
528 n 1)
529 if [ -n "$ts_iface" ]; then
530     echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
531     # DO NOT delete the physical default route! It crashes Colima.
532     # Use the /1 trick to mathematically override the default route.
533     sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
534     sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
535     sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
536     sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
537     enable_killswitch
538 fi
539
540 # 6. Stop gateway Docker containers (default) / Force-recreate warp if unhealthy (post-wake)
541 if [ "$POST_WAKE" = "true" ]; then
542     # 6a. Check for Roaming Network Mismatch (Bridged vs Enterprise Wi-Fi)
543     if [ -f "/tmp/nullexit_colima_mode.txt" ] && [ -f "$SCRIPT_DIR/../lan_p2p_detected" ]; then
544         current_colima_mode=$(cat /tmp/nullexit_colima_mode.txt)
545         p2p_allowed=$(cat "$SCRIPT_DIR/../lan_p2p_detected")
546         required_colima_mode="bridged"
547         [ "$p2p_allowed" == "false" ] && required_colima_mode="shared"
548
549         if [ "$current_colima_mode" != "$required_colima_mode" ]; then
550             warn "Network roaming mismatch detected!"
551             warn "Colima is running in '$current_colima_mode' mode but new network requires '$required_colima_mode'."
552             warn "Executing full gateway restart to protect against MAC-spoofing deauthentication..."
553             echo "[$(date +%Y-%m-%d %H:%M:%S)] [Roam] Mismatch ($current_colima_mode vs $required_colima_mode)
554             ). Triggering toggle.sh --restart." >> output.log
555             nohup bash "$SCRIPT_DIR/../toggle.sh" --restart >/dev/null 2>&1 &
556             exit 0
557         fi
558     fi
559
560     step "Checking Colima VM state"
561     colima_status=$(colima status 2>&1)
562     if echo "$colima_status" | grep -qi "running"; then
563         ok "Colima VM is running"
564         echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] Colima status check passed." >> output.log
565     else
566         warn "Colima VM is NOT running or unresponsive"
567         echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] Colima is unhealthy or dead! Output:
568         $colima_status" >> output.log
569         echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] Attempting to restart Colima..." >> output.log
570         colima restart >> output.log 2>&1 || warn "Failed to restart Colima"
571     fi
572
573     step "Checking warp container health (gluetun UDP tunnel)"
574     echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] Inspecting warp container state..." >> output.log
575     # The warp container is the most failure-prone link in the chain at wake/roam:
576     # its UDP NAT binding to Cloudflare's edge (162.159.192.1:2408) silently dies
577     # during a Wi-Fi roam. We verify the tunnel is actually passing traffic via curl.
578     WARP_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
579     echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] warp container status: $WARP_STATE" >> output.log
580
581     if [ "$WARP_STATE" = "missing" ]; then
582         warn "warp container is missing leaving gateway containers alone (full relaunch requires \
583         $SCRIPT_DIR/toggle.sh\)"
584         echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] warp container missing from Docker engine.
585         Requires full toggle.sh launch." >> output.log
586     else
587         tmp_err=$(mktemp)
588         if docker compose exec -T warp wget -qO- --timeout=3 https://www.cloudflare.com/cdn-cgi/trace 2>
589         $tmp_err | grep -q "warp=on"; then
590             ok "warp container is healthy and actively tunneling traffic (no recreate needed)"
591             echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] warp container passed live traffic healthcheck
592             (Cloudflare trace successful)." >> output.log
593         else
594             warn "warp container tunnel is dead (Docker status: $WARP_STATE) force-recreating all gateway
595             containers to restore network namespace"
596             echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] warp container is DEAD or STUCK. Forcing
597             recreation of gateway stack..." >> output.log
598             if [ -s "$tmp_err" ]; then
599                 err_msg=$(cat "$tmp_err" | tr -d '\r')
600                 echo "[$(date +%Y-%m-%d %H:%M:%S)] [Docker/Colima] Healthcheck failure details: $err_msg" >>
601                 output.log
602             fi
603             docker compose up -d --force-recreate >> output.log 2>&1 || warn "force-recreate failed"
604         fi
605     fi
606 fi

```

```

596     NEW_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing
597 ")
598     ok "warp container health after recreate: $NEW_STATE"
599     echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container health after recreation:
600 $NEW_STATE" >> output.log
601 fi
602 rm -f "$tmp_err"
603 fi
604 else
605     step "Stopping gateway Docker containers"
606     if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
607         if [ -f "docker-compose.yml" ]; then
608             docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were
609 running"
610         else
611             warn "docker-compose.yml not found in current directory"
612         fi
613     else
614         warn "Docker not available skipping"
615     fi
616 fi
617
618 # 6b. Stopping sleep prevention (caffeinate) default only
619 if [ "$POST_WAKE" = "false" ]; then
620     step "Stopping sleep prevention"
621     stop_pidfile_daemon "$PID_CAFFEINATE" "system sleep prevention"
622 else
623     step "Sleep prevention: leaving untouched (caffeinate already re-acquired by parent shell)"
624 fi
625
626 # 6c. Stopping DNS Watcher default only
627 if [ "$POST_WAKE" = "false" ]; then
628     step "Stopping DNS Watcher"
629     stop_pidfile_daemon "$PID_DNS_WATCHER" "background DNS Watcher"
630 else
631     step "DNS Watcher: leaving untouched (it keeps re-hijacking DNS on every roam)"
632 fi
633
634 # 7. Power-cycle Wi-Fi default only
635 if [ "$POST_WAKE" = "false" ]; then
636     step "Power-cycling Wi-Fi"
637     bounce_wifi_interfaces
638 else
639     step "Wi-Fi: leaving untouched (post-wake keeps the current Wi-Fi association)"
640 fi
641
642 # 8. Verify (default checks direct internet; post-wake verifies the gateway)
643 if [ "$POST_WAKE" = "true" ]; then
644     step "Verifying gateway is healthy after post-wake refresh"
645
646     # Wait a moment for tailscale + warp healthcheck to settle
647     sleep 3
648
649     GATEWAY_OK=false
650     if command -v dig >> output.log 2>&1; then
651         # Source procedural ports if available
652         if [ -f "/tmp/nullexit-ports.env" ]; then
653             source "/tmp/nullexit-ports.env"
654         fi
655         DNS_PORT=${DNS_PROXY_PORT:-5354}
656
657         # AdGuard is exposed at localhost:${DNS_PORT}. If this resolves, the warp tunnel
658         # is functioning properly because AdGuard routes upstream to warp.
659         if dig +tcp @127.0.0.1 -p "$DNS_PORT" google.com +short +timeout=5 >> output.log 2>&1; then
660             ok "Gateway DNS (AdGuard via localhost:${DNS_PORT}) works"
661             GATEWAY_OK=true
662         else
663             warn "Gateway DNS not responding yet"
664         fi
665     fi
666
667     # Direct host gateway check via Tailscale ping (relayed pong counts as success)
668     if command -v tailscale >> output.log 2>&1; then
669         TS_IP=""
670         [ -f .gateway_ip ] && TS_IP=$(read_adguard_ip || true)
671         if [ -n "$TS_IP" ] && tailscale ping --until-direct=false -c 1 --timeout 4s "$TS_IP" 2>> output.log |
672             grep -q pong; then
673             ok "Gateway $TS_IP reachable via Tailscale"
674             GATEWAY_OK=true
675         else
676             warn "Tailscale ping to gateway did not pong exit-node may still be re-establishing"
677         fi
678     fi
679 fi

```

```

673     fi
674 fi
675
676 if [ "$GATEWAY_OK" = "true" ]; then
677     echo -e "  ${GREEN}${BOLD} Gateway is healthy after post-wake refresh.${NC}"
678 else
679     echo -e "  ${YELLOW}${BOLD} Gateway recovery is in-progress.${NC}"
680     echo "    The route may need another 30s before queries succeed."
681     echo "    If it stays broken for >60s, run: bash \"\${SCRIPT_DIR}/toggle.sh\""
682 fi
683 else
684     verify_internet_connectivity
685 fi
686
687 # Summary
688 echo ""
689 echo -e "${BOLD}${NC}"
690 if [ "$POST_WAKE" = "true" ]; then
691     echo -e "${BOLD}POST-WAKE SUMMARY${NC}"
692 else
693     echo -e "${BOLD}RECOVERY SUMMARY${NC}"
694 fi
695 echo -e "${BOLD}${NC}"
696
697 # Show final DNS state
698 FINAL_DNS=$(networksetup -getdnsservers "Wi-Fi" 2>/dev/null || echo "N/A")
699 echo "  DNS (Wi-Fi):  $FINAL_DNS"
700
701 FINAL_DNS2=$(networksetup -getdnsservers "Ethernet" 2>/dev/null || true)
702 if [[ -z "$FINAL_DNS2" || "$FINAL_DNS2" == *"not a recognized"* || "$FINAL_DNS2" == *"Error"* ]]; then
703     FINAL_DNS2="N/A (Not configured)"
704 fi
705 echo "  DNS (Ethernet): $FINAL_DNS2"
706
707 echo ""
708
709 if [ "$POST_WAKE" = "false" ]; then
710     if [ "$INTERNET_OK" = "true" ]; then
711         echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
712     else
713         echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
714         echo "    This might mean:"
715         echo "    - Your Wi-Fi is disconnected from the router"
716         echo "    - Tailscale is still interfering (try restarting tailscaled:"
717         echo "      brew services restart tailscale)"
718         echo "    - You need to re-connect to Wi-Fi in the menu bar"
719         echo "    - Try toggling Wi-Fi off/on in the menu bar"
720         echo ""
721         echo "    Or try manually:"
722         echo "    networksetup -setdnsservers Wi-Fi empty"
723         echo "    tailscale down"
724         echo "    sudo route delete -net 0.0.0.0/1"
725         echo "    sudo route delete -net 128.0.0.0/1"
726         echo "    brew services restart tailscale"
727     fi
728 fi
729
730 # Reset sharing services to prevent AirDrop freezes after interface changes
731 # This is the SAME step in BOTH modes (post-wake inherits the previous fix from 10.27).
732 echo -e "  Resetting macOS sharing services (AirDrop/AirPlay)..."
733 reset_sharing_services
734
735 echo ""
736 if [ "$POST_WAKE" = "true" ]; then
737     echo -e "${YELLOW}${BOLD}Post-wake refresh complete.${NC}"
738     # Remove the WARP Watcher inhibit marker now that the gateway is fully healthy.
739     # The watcher will resume normal liveness monitoring on its next 5s poll.
740     rm -f "$WARP_INHIBIT_FILE"
741     echo "[$(date -u +%FT%TZ)] POST-WAKE complete  WARP Watcher inhibit released." >> output.log
742 else
743     echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
744 fi
745 echo ""

```

8 setup.sh

```

1 #!/bin/bash
2 # setup.sh nullexit one-shot setup script (macOS)
3 # Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
4 set -euo pipefail
5
6 # Source common formatting and helper functions
7 source "$(dirname "$0")/scripts/common.sh"
8
9 # Run common installation logic
10 source "$(dirname "$0")/scripts/setup-common.sh"
11
12 # Compile macOS Shortcuts
13 echo -e "\n${BOLD}Compiling macOS Desktop Shortcuts...${NC}"
14 if command -v osacompile &> /dev/null; then
15     osacompile -o "Toggle Gateway.app" "Toggle-Gateway.applescript" >/dev/null 2>&1
16     osacompile -o "Recover Gateway.app" "Recover-Gateway.applescript" >/dev/null 2>&1
17     echo "    Created 'Toggle Gateway.app' and 'Recover Gateway.app'"
18 fi
19
20 # Done
21 echo ""
22 echo -e "${GREEN}${BOLD}${NC}"
23 echo -e "${GREEN}${BOLD}                Setup complete!                ${NC}"
24 echo -e "${GREEN}${BOLD}${NC}"
25 echo ""
26 echo -e "${BOLD}One manual step remaining approve the exit node:${NC}"
27 echo "    https://login.tailscale.com/admin/machines"
28 echo "    Find your gateway '...' 'Edit route settings' enable Exit Node"
29 echo ""
30
31 echo -e "${YELLOW}${BOLD}Sudoers / Background Execution Requirements:${NC}"
32 echo "    To allow silent, non-interactive execution of the firewall and network routes (especially during
33 sleep/wake recovery), you must configure passwordless sudo."
34 echo "    Run this exact command in your terminal (grants are scoped to least privilege):"
35 echo "    echo \"\$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /sbin/route, /sbin/
36 ifconfig, /usr/bin/dscacheutil -flushcache, /usr/bin/killall -HUP mDNSResponder, /usr/bin/killall
37 sharingd, rapportd, /usr/bin/pkill -9 tailscaled, /usr/bin/pkill -f dns-proxy.py, /bin/kill, /usr/bin
38 /true, /opt/homebrew/bin/brew services * tailscale, /opt/homebrew/bin/brew uninstall tailscale, /usr
39 /local/bin/brew services * tailscale, /usr/local/bin/brew uninstall tailscale, /usr/bin/python3 -I \
40 \$PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I \$PWD/scripts/dns-proxy.py, /usr/local/bin/
41 python3 -I \$PWD/scripts/dns-proxy.py\" | sudo tee /etc/sudoers.d/nullexit >/dev/null && sudo visudo
42 -cf /etc/sudoers.d/nullexit"
43 echo ""
44
45 if [[ -n "${TS_IP:-}" ]]; then
46     echo -e "${BOLD}AdGuard Home dashboard:${NC}    http://${TS_IP}:3000    (username: admin)"
47     echo ""
48 fi
49
50 echo -e "${BOLD}To toggle the gateway on/off:${NC}"
51 echo "    macOS: double-click the 'Toggle Gateway' app icon!"
52 echo "    Windows: double-click Toggle-Gateway.bat"
53 echo ""
54
55 if [[ "$OSTYPE" == "darwin"* ]]; then
56     # LuLu localhost filtering check
57     if [ -d "/Applications/LuLu.app" ] || systemextensionsctl list 2>/dev/null | grep -q "com.objective-
58 see.lulu"; then
59         echo -e "${YELLOW}${BOLD}LuLu Firewall Detected:${NC}"
60         echo "    LuLu's localhost (loopback) filtering is OFF by default. To properly protect the"
61         echo "    local nullexit proxy ports from malicious local processes, you must manually enable"
62         echo "    'Allow Loopback' (or block it selectively) in LuLu's Preferences -> Rules."
63         echo ""
64     fi
65
66     echo -e "${YELLOW}${BOLD}Note on macOS Tailscale Sandboxing:${NC}"
67     echo "    The Tailscale GUI app (tailscale-app) installed on macOS is sandboxed."
68     echo "    While you cannot run a Tailscale SSH server on this Mac host, you can"
69     echo "    still SSH outbound to other devices on the mesh using their MagicDNS"
70     echo "    addresses directly (e.g. ssh user@device.tailnet-name.ts.net)."
71     echo ""
72 fi

```

9 scripts/setup-linux.sh

```

1 #!/bin/bash

```

```

2 # scripts/setup-linux.sh nullexit one-shot setup script (Linux)
3 # Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
4 set -euo pipefail
5
6 # Source common formatting and helper functions
7 source "${dirname "$0"}/common.sh"
8
9 # Run common installation logic
10 source "${dirname "$0"}/setup-common.sh"
11
12 # Compile Linux Desktop Shortcuts
13 echo -e "\n${BOLD}Creating Linux Desktop Shortcuts...${NC}"
14 DIR="$(pwd)"
15 cat <<EOF > "Toggle Gateway.desktop"
16 [Desktop Entry]
17 Version=1.0
18 Name=Toggle Gateway
19 Comment=Start or Stop Nullexit Gateway
20 Exec=bash -c "cd '$DIR' && ./toggle.sh; read -p 'Press Enter to close...' dummy"
21 Icon=utilities-terminal
22 Terminal=true
23 Type=Application
24 Categories=Network;Security;
25 EOF
26
27 cat <<EOF > "Recover Gateway.desktop"
28 [Desktop Entry]
29 Version=1.0
30 Name=Recover Gateway
31 Comment=Emergency DNS and Network Recovery
32 Exec=bash -c "cd '$DIR' && ./recover.sh; read -p 'Press Enter to close...' dummy"
33 Icon=utilities-terminal
34 Terminal=true
35 Type=Application
36 Categories=Network;Security;
37 EOF
38 echo "    Created 'Toggle Gateway.desktop' and 'Recover Gateway.desktop'"
39
40 # Done
41 echo ""
42 echo -e "${GREEN}${BOLD}${NC}"
43 echo -e "${GREEN}${BOLD}                Setup complete!                ${NC}"
44 echo -e "${GREEN}${BOLD}${NC}"
45 echo ""
46 echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
47 echo "    https://login.tailscale.com/admin/machines"
48 echo "    Find your gateway '...' 'Edit route settings' enable Exit Node"
49 echo ""
50
51 if [[ -n "${TS_IP:-}" ]]; then
52     echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
53     echo ""
54 fi
55
56 echo -e "${BOLD}To toggle the gateway on/off:${NC}"
57 echo "    Linux: double-click the 'Toggle Gateway' desktop icon!"
58 echo "    (or run ./toggle.sh in terminal)"
59 echo ""

```

10 docker-compose.yml

```

1 services:
2   warp:
3     image: qmcgaw/gluetun:v3.41.1
4     container_name: warp
5     hostname: nullexit-gateway
6     cap_add:
7       - NET_ADMIN
8     devices:
9       - /dev/net/tun:/dev/net/tun
10    ports:
11      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/udp" # AdGuard DNS: host binds 5354 (the mapped DNS-
12        proxy port); 5335 is AdGuard's in-container port
13      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/tcp" # AdGuard DNS (TCP)
14      - "41642:41642/udp" # Tailscale WireGuard (direct connection to host)
15      - "127.0.0.1:${SOCKS_PROXY_PORT:-1080}:1080/tcp" # SOCKS5 proxy (routed through WARP tunnel)

```



```

15     - "127.0.0.1:${TOR SOCKS_PORT:-9050}:9050/tcp" # Tor SOCKS5 proxy (procedurally generated)
16     - "127.0.0.1:${TOR_CONTROL_PORT:-9051}:9051/tcp" # Tor Control port (procedurally generated)
17 environment:
18     - TZ=America/New_York
19     - VPN_SERVICE_PROVIDER=custom
20     - VPN_TYPE=wireguard
21     - WIREGUARD_PRIVATE_KEY=${WIREGUARD_PRIVATE_KEY}
22     - WIREGUARD_PUBLIC_KEY=${WIREGUARD_PUBLIC_KEY}
23     - WIREGUARD_ADDRESSES=${WIREGUARD_ADDRESSES}
24     - WIREGUARD_ENDPOINT_IP=${WARP_ENDPOINT_1:-162.159.192.1}
25     - WIREGUARD_ENDPOINT_PORT=2408
26     # Set to 25s (WireGuard default) to beat standard 30s hardware NAT/firewall UDP timeouts
27     - WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL=25s
28     # MTU Fix (Critical): Prevent packet fragmentation from double-encryption
29     - WIREGUARD_MTU=1280
30     # Give WARP more time to establish the connection before internal restart (default is 6s which is
    too short)
31     - HEALTH_VPN_DURATION_INITIAL=30s
32     # Use plaintext DNS resolvers instead of DNS-over-TLS (DoT) to prevent startup failures on networks
    where port 853 is blocked
33     - DNS_UPSTREAM_RESOLVER_TYPE=plain
34     - DNS_UPSTREAM_PLAIN_ADDRESSES=1.1.1.1:53,8.8.8.8:53
35     # Allow local network traffic to bypass WARP so Tailscale can establish fast, direct peer-to-peer
    connections (no DERP relays) on the LAN
36     - FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8
37     # Built-in DNS blocking for ads, malicious domains, and tracking
38     # NOTE: This is the massive "Hagezi-style" built-in blocklist.
39     # You can turn these to 'on' for maximum security if you have spare RAM
40     # or are running this on a cloud exit node with adequate autonomous resources (requires approx 1-2
    GB of RAM).
41     # Disabled intentionally: AdGuard Home provides the DNS filtering layer instead.
42     # If AdGuard fails, there is no fallback filtering, but the healthcheck
43     # dependency chain ensures Tailscale won't come up if AdGuard is down.
44     - BLOCK_MALICIOUS=off
45     - BLOCK_SURVEILLANCE=off
46     - BLOCK_ADS=off
47 dns:
48     - 1.1.1.1
49 sysctls:
50     - net.ipv4.ip_forward=1
51 volumes:
52     - ./post-rules.txt:/iptables/post-rules.txt:ro
53 restart: unless-stopped
54 healthcheck:
55     test: ["CMD-SHELL", "/gluetun-entrypoint healthcheck"]
56     interval: 2s
57     timeout: 5s
58     retries: 15
59     start_period: 60s
60 logging:
61     driver: "json-file"
62     options:
63         max-size: "10m"
64         max-file: "3"
65
66 tailscale:
67     image: tailscale/tailscale:v1.98.4
68     container_name: tailscale
69     network_mode: service:warp
70     depends_on:
71         warp:
72             condition: service_healthy
73         adguardhome:
74             condition: service_healthy
75     cap_add:
76         - NET_ADMIN
77         - NET_RAW
78         - SYS_MODULE
79     sysctls:
80         - net.ipv4.ip_forward=1
81         - net.ipv6.conf.all.forwarding=1
82     environment:
83         - TZ=America/New_York
84         - TS_EXTRA_ARGS=--advertise-exit-node
85         - PORT=41642
86         - TS_USERSPACE=false
87         - TS_AUTHKEY=${TS_AUTHKEY}
88         - TS_STATE_DIR=/var/lib/tailscale
89         # (Note: For headless/cloud deployments, you can bypass this footgun by
90         # using Tailscale Ephemeral Keys which auto-clean themselves).
91         - TS_AUTH_ONCE=true

```

```

92 # MTU Optimization: Force Tailscale packets to be smaller than WARP's 1280 MTU
93 # to prevent WireGuard-in-WireGuard fragmentation, completely eliminating bufferbloat
94 - TS_DEBUG_MTU=1200
95 # Enable Tailscale's built-in SOCKS5 proxy to route out of the WARP tunnel
96 - TS_SOCKS5_SERVER=:1080
97 devices:
98   - /dev/net/tun:/dev/net/tun
99 volumes:
100   - tailscale_state:/var/lib/tailscale
101 restart: unless-stopped
102 logging:
103   driver: "json-file"
104   options:
105     max-size: "10m"
106     max-file: "3"
107
108 routing-fix:
109   image: nullexit-routing-fix:v1.0.0
110   build:
111     context: ./scripts
112     dockerfile: Dockerfile.routing-fix
113   container_name: routing-fix
114   network_mode: service:warp
115   # NOTE: Deliberately NOT sharing tailscale's PID namespace the shared
116   # namespace caused routing-fix to be killed whenever tailscale restarted
117   # (e.g. during gluetun health-check flaps), which dropped all routing
118   # table changes and left traffic leaking through eth0 instead of tun0.
119   cap_add:
120     - NET_ADMIN
121     - SYSLOG
122   depends_on:
123     warp:
124       condition: service_healthy
125     # rule-compiler dependency removed: it is now a hash-gated one-off run by
126     # toggle.sh before 'up' (15.8.8). This service boots on the existing
127     # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
128   environment:
129     - TZ=America/New_York
130     - WARP_ENDPOINT=${WARP_ENDPOINT_1:-162.159.192.1}
131     - BLOCKED_COUNTRIES=${BLOCKED_COUNTRIES}
132     - TAILSCALE_ALLOW_LAN_P2P=${TAILSCALE_ALLOW_LAN_P2P:-false}
133   volumes:
134     - ./scripts/routing-fix.sh:/routing-fix.sh:ro
135     # SECURITY: ./adguard/work/ is bind-mounted into this container.
136     # Never write to it from the host or another container. The
137     # single allowed writer is 'rule-compiler' container manual
138     # edits corrupt the AdGuard cache + BoltDB session store.
139     - ./adguard/work/userfilters:/userfilters:ro
140     - ./adguard/work:/adguard_work:ro
141     - ./app
142 restart: unless-stopped
143 stop_grace_period: 1s
144 logging:
145   driver: "json-file"
146   options:
147     max-size: "1m"
148     max-file: "1"
149
150 adguardhome:
151   # WARNING: AdGuard is pre-configured with hardcoded credentials (admin/nullexit).
152   # This is fine for a local Mac running behind a NAT, but if you are deploying
153   # this on a public cloud VPS with port 80/3000 exposed, change this immediately!
154   image: adguard/adguardhome:v0.107.77
155   container_name: adguardhome
156   network_mode: service:warp
157   depends_on:
158     warp:
159       condition: service_healthy
160     # rule-compiler dependency removed: it is now a hash-gated one-off run by
161     # toggle.sh before 'up' (15.8.8). This service boots on the existing
162     # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
163   environment:
164     - TZ=America/New_York
165   healthcheck:
166     test: ["CMD-SHELL", "nc -z 127.0.0.1 5335 || exit 1"]
167     interval: 2s
168     timeout: 5s
169     retries: 10
170   volumes:
171     # SECURITY: ./adguard/work/ is bind-mounted into this container.
172     # Never write to it from the host or another container. The

```

```

173 # single allowed writer is 'rule-compiler' container manual
174 # edits corrupt the AdGuard cache + BoltDB session store.
175 - ./adguard/work:/opt/adguardhome/work
176 - ./adguard/conf:/opt/adguardhome/conf
177 restart: unless-stopped
178 stop_grace_period: 2s
179 logging:
180   driver: "json-file"
181   options:
182     max-size: "10m"
183     max-file: "3"
184
185 rule-compiler:
186   image: nullexit-rule-compiler:v1.0.0
187   build:
188     context: ./scripts
189     dockerfile: Dockerfile.rule-compiler
190   container_name: rule-compiler
191   # 'compile' profile: excluded from 'docker compose up'. toggle.sh runs it
192   # on-demand via 'docker compose run --build --rm rule-compiler', hash-gated on
193   # black_list.txt/white_list.txt, so ~340k rules are re-deduped only when a list
194   # actually changes instead of every toggle (~28s saved on the common path, 15.8.8).
195   profiles: ["compile"]
196   volumes:
197     - ../app
198
199 tor:
200   build: ./docker/tor
201   image: nullexit-tor:v1.0.0
202   container_name: tor
203   network_mode: service:warp
204   depends_on:
205     warp:
206       condition: service_healthy
207   environment:
208     - TZ=America/New_York
209     - TOR_USE_BRIDGES=${TOR_USE_BRIDGES:-false}
210     - TOR_BRIDGE_LINES_FILE=${TOR_BRIDGE_LINES_FILE:-./tor-bridges.txt}
211     - TOR_PASSWORD=${ADGUARD_PASSWORD:-nullexit}
212     - TOR_EXCLUDE_EXIT_NODES=${TOR_EXCLUDE_EXIT_NODES:-{us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il}}
213   working_dir: /nullexit
214   volumes:
215     - ../:/nullexit:ro
216     - ./output.log:/output.log:rw
217   restart: unless-stopped
218   logging:
219     driver: "json-file"
220     options:
221       max-size: "10m"
222       max-file: "3"
223
224 volumes:
225   tailscale_state:
226
227 networks:
228   default:
229     ipam:
230       config:
231         - subnet: 10.200.1.0/24

```

11 scripts/common.sh

```

1 #!/bin/bash
2 # common.sh shared bash utilities for nullexit scripts
3
4 # Formatting
5 RED='\033[0;31m'
6 GREEN='\033[0;32m'
7 YELLOW='\033[1;33m'
8 BLUE='\033[0;34m'
9 BOLD='\033[1m'
10 NC='\033[0m'
11
12 # Constants & Environment
13 export MARKER_FILE="/tmp/nullexit-gateway-active.marker"
14 export PID_CAFFEINATE="/tmp/nullexit-caffeinate.pid"

```

```

15 export PID_DNS_WATCHER="/tmp/nullexit-dns-watcher.pid"
16 export DEFAULT_WARP_ENDPOINT_1="162.159.192.1"
17 export DEFAULT_WARP_ENDPOINT_2="162.159.193.1"
18
19 setup_standard_path() {
20     export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
21 }
22
23 step() { echo -e "\n[ $(date '+%Y-%m-%d %H:%M:%S') ] ${BLUE}${BOLD} ${NC}"; }
24 ok() { echo -e "[ $(date '+%Y-%m-%d %H:%M:%S') ] ${GREEN} ${NC}"; }
25 warn() { echo -e "[ $(date '+%Y-%m-%d %H:%M:%S') ] ${YELLOW} ${NC}"; }
26 fail() { echo -e "[ $(date '+%Y-%m-%d %H:%M:%S') ] ${RED} ${NC}"; }
27 die() { echo -e "\n[ $(date '+%Y-%m-%d %H:%M:%S') ] ${RED} ${NC}\n"; exit 1; }
28
29 # Helper Functions
30
31 # Append a timestamped line to output.log (best-effort; never fails the caller).
32 # Used to record lifecycle breadcrumbs (EXEC/EXIT markers, phase changes) so a
33 # START/STOP that is killed or window-closed still leaves a retraceable trail.
34 log_line() {
35     echo "[ $(date '+%Y-%m-%d %H:%M:%S') ] *" >> "${LOG_FILE:-output.log}" 2>/dev/null || true
36 }
37
38 # Run a lifecycle command with FULL observability: log the command, stream its
39 # combined stdout+stderr to the terminal AND append it to output.log via tee,
40 # then log the exit code explicitly. Returns the command's real exit code (not
41 # tee's). This closes the blind spot where 'docker compose up' / 'colima start'
42 # printed only to the terminal so when they fail (or the run is interrupted),
43 # output.log shows exactly what the last command was and whether it succeeded.
44 #
45 # Usage: run_logged docker compose up -d --build
46 run_logged() {
47     local logf="${LOG_FILE:-output.log}"
48     log_line "EXEC: *"
49     # PIPESTATUS[0] is the command's exit code; tee (PIPESTATUS[1]) is ignored.
50     "$@" 2>&1 | tee -a "$logf"
51     local rc=${PIPESTATUS[0]}
52     log_line "EXIT $rc: *"
53     return "$rc"
54 }
55
56 # Pure-bash timeout (no dependency on GNU coreutils' timeout)
57 # Runs a command with a safety cutoff so a wedged daemon can't hang the script.
58 run_with_timeout() {
59     local timeout_sec="$1"
60     shift
61     if [ $# -eq 0 ]; then return 1; fi
62
63     "$@" &
64     local cmd_pid=$!
65     CURRENT_BG_PID="$cmd_pid"
66
67     (
68         sleep "$timeout_sec"
69         if kill -0 "$cmd_pid" 2>> output.log; then
70             echo -e "\n[Timeout] Command '*' exceeded $timeout_sec seconds. Terminating..." >&2
71             kill -15 "$cmd_pid" 2>> output.log || true
72             sleep 2
73             if kill -0 "$cmd_pid" 2>> output.log; then
74                 kill -9 "$cmd_pid" 2>> output.log || true
75             fi
76         fi
77     ) &
78     local watcher_pid=$!
79
80     # Disable set -e temporarily to capture exit status of wait
81     set +e
82     wait "$cmd_pid" 2>> output.log
83     local exit_status=$?
84     set -e
85
86     # Kill the watcher since the command finished
87     kill "$watcher_pid" 2>> output.log && wait "$watcher_pid" 2>> output.log || true
88
89     CURRENT_BG_PID=""
90     return "$exit_status"
91 }
92
93 # Start Colima and return the moment Docker is actually reachable, instead of
94 # blocking on colima 0.10.3's post-provision hang. Observed on macOS/vz: the VM
95 # reaches READY and creates the 'colima' docker context in ~10-15s, then

```

```

96 # 'colima start' frequently fails to return for ~110s (until a watchdog kills it)
97 # even though Docker is fully usable the entire time. We poll the colima docker
98 # context and proceed as soon as it answers, then reap the (usually hung) starter.
99 # The VM keeps running and 'colima status' stays accurate. $@ = colima start args.
100 # Returns 0 once Docker responds, 1 if it never does within the cap.
101 colima_start_until_ready() {
102     local max_wait=90 waited=0 ready=false start_pid
103     echo "[$(date +%Y-%m-%d %H:%M:%S')] EXEC(bg): colima start $*" >> output.log
104     colima start "$@" >> output.log 2>&1 &
105     start_pid=$!
106     while [ "$waited" -lt "$max_wait" ]; do
107         if run_with_timeout 5 docker --context colima info >/dev/null 2>&1; then
108             ready=true
109             break
110         fi
111         # If colima start exited on its own (genuine finish or hard failure), stop waiting.
112         if ! kill -0 "$start_pid" 2>/dev/null; then
113             run_with_timeout 5 docker --context colima info >/dev/null 2>&1 && ready=true
114             break
115         fi
116         sleep 3
117         waited=$((waited + 3))
118     done
119     # Reap the starter (harmless if it already exited); this does NOT stop the VM.
120     kill "$start_pid" 2>/dev/null || true
121     wait "$start_pid" 2>/dev/null || true
122     if [ "$ready" = "true" ]; then
123         docker context use colima >/dev/null 2>&1 || true
124         echo "[$(date +%Y-%m-%d %H:%M:%S')] Colima Docker ready after ~${waited}s (starter reaped)." >>
            output.log
125         return 0
126     fi
127     echo "[$(date +%Y-%m-%d %H:%M:%S')] Colima Docker NOT ready after ${max_wait}s." >> output.log
128     return 1
129 }
130
131 # Hash the files that feed the long-running locally-built images that 'docker
132 # compose up' starts (routing-fix, tor). toggle.sh uses this to skip
133 # 'docker compose --build' when nothing changed: BuildKit re-evaluating those
134 # images on the 600MB Colima VM costs ~30s per toggle even when every layer is
135 # CACHED (15.8.8). Includes the Dockerfiles, everything they COPY (routing-fix.sh,
136 # tor/entrypoint.sh, the logger/ Go source), and docker-compose.yml. The
137 # rule-compiler image is NOT here it is a 'compile'-profiled one-off, built+run
138 # by the gated compile step (15.8.8). Prints a sha256 hex digest, or empty.
139 compute_build_inputs_hash() {
140     local root="${SCRIPT_DIR:-.}" f
141     {
142         for f in \
143             "$root/scripts/Dockerfile.routing-fix" \
144             "$root/scripts/routing-fix.sh" \
145             "$root/docker/tor/Dockerfile" \
146             "$root/docker/tor/entrypoint.sh" \
147             "$root/docker-compose.yml"; do
148             printf ':%s:.' "$f"; cat "$f" 2>/dev/null
149         done
150         find "$root/scripts/logger" -type f 2>/dev/null \
151             | LC_ALL=C sort | while IFS= read -r f; do printf ':%s:.' "$f"; cat "$f" 2>/dev/null; done
152     } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
153 }
154
155 # Hash the ad-block lists the user controls: black_list.txt (custom blocks) and
156 # white_list.txt (the allow-list that governs what DNS resolves through). toggle.sh
157 # gates the ~28s rule recompile on this editing either list changes the digest
158 # and triggers a one-off recompile of compiled_rules.txt + ip_blocklist.ipset;
159 # an unchanged toggle skips it entirely (15.8.8). Prints a sha256 hex digest.
160 compute_rules_hash() {
161     local root="${SCRIPT_DIR:-.}"
162     { printf ':black:.'; cat "$root/black_list.txt" 2>/dev/null
163       printf ':white:.'; cat "$root/white_list.txt" 2>/dev/null
164     } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
165 }
166
167 # Check if the gateway is active (either containers are running, or host DNS is hijacked)
168 is_gateway_active() {
169     # 1. Check if containers are running (suppress stderr to avoid errors if docker is down)
170     if run_with_timeout 15 docker compose ps --status running 2>/dev/null | grep -q 'warp'; then
171         echo "[$(date +%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (warp container is running)" >> "
            $LOG_FILE"
172         return 0
173     fi
174 }

```

```

175 # 2. Check if host DNS was hijacked (not 1.1.1.1 and not default/empty)
176 local current_dns=""
177 local active_svc=""
178 if [[ "$OSTYPE" == "darwin"* ]]; then
179     active_svc=$(get_active_service)
180     current_dns=$(networksetup -getdnsservers "$active_svc" 2>> output.log || true)
181 else
182     current_dns=$(resolvectl dns 2>/dev/null | awk '/Global|Link/ {for(i=4;i<=NF;i++) print $i}' | head -
183         n1)
184 fi
185 if [[ -n "$current_dns" && "$current_dns" != "1.1.1.1" && ! "$current_dns" =~ "There aren't any DNS
186     Servers" ]]; then
187     echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (DNS was hijacked: $current_dns)" >> "$
188         LOG_FILE"
189     return 0
190 fi
191 # 3. Check if SOCKS proxy is enabled (macOS only)
192 if [[ "$OSTYPE" == "darwin"* ]]; then
193     local socks_proxy
194     socks_proxy=$(networksetup -getsocksfirewallproxy "$active_svc" 2>> output.log || true)
195     if echo "$socks_proxy" | grep -q "Enabled: Yes"; then
196         echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (SOCKS proxy enabled)" >> "$LOG_FILE"
197         return 0
198     fi
199 fi
200 echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: FALSE (Containers down, DNS clean, SOCKS
201     disabled)" >> "$LOG_FILE"
202 return 1
203 }
204 # Decide whether the gateway SHOULD be running, from persisted *intent* rather
205 # than a live health probe. This is the correct signal for toggle.sh's
206 # STOP-vs-START decision: a disable should tear down whatever the last START
207 # left behind it should NOT depend on whether Docker/Colima happens to be
208 # reachable right now.
209 #
210 # Why this exists: using is_gateway_active() for the decision means every toggle
211 # (including *disable*) runs 'docker compose ps' first. When Colima is down the
212 # socket is absent, that call blocks for the full run_with_timeout window, and
213 # under 'set -e' + the ERR trap a non-zero result at the 'if then else fi'
214 # boundary could abort the script BEFORE the START body runs so the very path
215 # that would boot Colima never executed (observed as 'EXIT: code=1 phase=start').
216 #
217 # Signal source: MARKER_FILE is written at the end of a successful START
218 # (write_gateway_active_marker) and cleared at the top of STOP and in
219 # cleanup_handler. Reading it is instant and cannot hang or abort.
220 #
221 # ECHOES "running" or "stopped" to stdout (and ALWAYS returns 0). Callers read
222 # the string: [ "$(gateway_state)" = "running" ] &&
223 #
224 # CRITICAL why this echoes instead of using a 0/1 return code: on macOS
225 # /bin/bash 3.2, with 'set -e' + an 'ERR' trap + 'set -x' (DEBUG_TRACE=true) all
226 # active, a function that does 'return 1' triggers a SPURIOUS 'set -e' abort in
227 # the caller even when the call is guarded ('if f; then', or even 'f || rc=$?').
228 # The 'return N' from the function, traced by 'set -x', trips the ERR trap. This
229 # is a genuine bash-3.2 interaction (see devref 15.11.8). Echoing a result and
230 # always returning 0 sidesteps 'return'/'set -e'/'set -x' entirely the function
231 # can never contribute a non-zero status, so no caller can be aborted by it.
232 gateway_state() {
233     # Primary: persisted intent. Marker present => a START completed and no STOP
234     # has cleared it yet.
235     #
236     # NOTE: toggle.sh runs a "defensive stale-marker clear" near the top of every
237     # invocation (before this decision) to stop the wake-watcher firing against a
238     # crashed gateway. That would wipe the marker before we read it, so toggle.sh
239     # captures the marker's existence into GATEWAY_MARKER_AT_STARTUP *before* that
240     # clear and we honor it here. If the variable is unset (other callers, or the
241     # --restart pre-check which runs before the defensive clear), fall back to the
242     # live marker file.
243     if [ -n "${GATEWAY_MARKER_AT_STARTUP:-}" ]; then
244         if [ "$GATEWAY_MARKER_AT_STARTUP" = "yes" ]; then echo "running"; return 0; fi
245     elif [ -f "$MARKER_FILE" ]; then
246         echo "running"; return 0
247     fi
248     # Marker absent. Normally this means "stopped". But cover the edge case where
249     # the marker was lost (e.g. /tmp cleared, or gateway started by an older build)
250     # yet containers are genuinely up reconcile via is_gateway_active, but ONLY

```



```

252 # when Docker is actually reachable, so a dead-socket probe can never hang. On
253 # macOS Docker runs in the Colima VM; if that socket is absent, the gateway
254 # definitively is not running.
255 local docker_up="no"
256 if [[ "$OSTYPE" == "darwin"* ]]; then
257     if [ -S /var/run/docker.sock ] || [ -S "$HOME/.colima/default/docker.sock" ]; then docker_up="yes";
258     fi
259 else
260     if [ -S /var/run/docker.sock ]; then docker_up="yes"; fi
261 fi
262 if [ "$docker_up" = "yes" ] && is_gateway_active; then
263     echo "running"; return 0
264 fi
265 echo "stopped"
266 return 0
267 }
268
269 # Reads an environment variable from a file (default: .env), strips quotes and whitespace
270 read_env_var() {
271     local var_name="$1"
272     local env_file="${2:-.env}"
273     grep -E "$var_name=" "$env_file" 2>/dev/null | cut -d'=' -f2- | tr -d "\"\`" | sed -e 's/^[[:space:]]*://' -e 's/[[:space:]]*$//' || echo ""
274 }
275
276 # Load KILL_SWITCH variable from .env (defaults to false if not found/empty)
277 export KILL_SWITCH
278 KILL_SWITCH=$(read_env_var "KILL_SWITCH" 2>/dev/null | tr '[:upper:]' '[:lower:]')
279 if [ -z "$KILL_SWITCH" ]; then
280     KILL_SWITCH="false"
281 fi
282
283 # Function to get the active network service name (e.g., "Wi-Fi" or "USB 10/100 LAN")
284 get_active_service() {
285     if [[ "$OSTYPE" == "darwin"* ]]; then
286         local iface
287         iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
288
289         # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical
290         # interface
291         if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
292             for i in en0 en1 en2 en3; do
293                 if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep
294                     -q "inet "; then
295                     iface="$i"
296                     break
297                 fi
298             done
299         fi
300
301         # Default fallback if still empty or not enX
302         if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
303             iface="en0"
304         fi
305
306         # Map interface (e.g., en0) to service name (e.g., Wi-Fi)
307         local service
308         service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B 1 "Device: $iface" | head -n
309             1 | sed -E 's/^\([0-9\*]+\)\s*//')
310
311         if [ -n "$service" ]; then
312             echo "$service"
313         else
314             echo "Wi-Fi"
315         fi
316     else
317         # Get the interface used for default internet routing
318         local iface
319         iface=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
320         if [ -z "$iface" ]; then
321             echo ""
322             return 1
323         fi
324         echo "$iface"
325     fi
326 }
327
328 get_en0_service() {
329     local service

```

```

327 service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B1 "Device: en0" | head -1 | sed
-E 's/^\([0-9\*]+\)/ /' || true)
328 if [ -z "$service" ]; then
329     echo "Wi-Fi"
330 else
331     echo "$service"
332 fi
333 }
334
335 get_warp_endpoint_1() { read_env_var "WARP_ENDPOINT_1" ".env" | grep -Eo '[0-9\.]+' || echo "
$DEFAULT_WARP_ENDPOINT_1"; }
336 get_warp_endpoint_2() { read_env_var "WARP_ENDPOINT_2" ".env" | grep -Eo '[0-9\.]+' || echo "
$DEFAULT_WARP_ENDPOINT_2"; }
337
338 restart_tailscaled_daemon() {
339     echo " [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
340     echo " [Tailscale Recovery] Attempting to restart tailscaled service..."
341
342     if run_with_timeout 15 brew services restart tailscale >> output.log 2>&1; then
343         echo " [Tailscale Recovery] Successfully restarted tailscaled (user service)."
344     elif run_with_timeout 15 sudo -n brew services restart tailscale >> output.log 2>&1; then
345         echo " [Tailscale Recovery] Successfully restarted tailscaled (system service)."
346     else
347         echo " [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed
."
348     fi
349     sleep 3
350 }
351
352 disconnect_tailscale_host() {
353     local ts_bin="${1:-tailscale}"
354     if command -v "$ts_bin" >> output.log 2>&1; then
355         echo " Disconnecting host Tailscale from mesh..."
356
357         # Check if actually connected first (with timeout tailscaled could be wedged)
358         local status_ok=false
359         if run_with_timeout 5 "$ts_bin" status >> output.log 2>&1; then
360             status_ok=true
361         else
362             local status_exit=$?
363             # If it was a quick exit code 1 (disconnected), it is still reachable
364             if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
365                 status_ok=true
366             fi
367         fi
368
369         if [ "$status_ok" = "false" ]; then
370             restart_tailscaled_daemon
371         fi
372
373         # Explicitly reset any exit-node so it doesn't linger.
374         "$ts_bin" up >> output.log 2>&1 || true
375         if run_with_timeout 10 "$ts_bin" set --ssh=true --accept-dns=false --exit-node= >> output.log 2>&1;
376         then
377             echo " [] Exit-node preference cleared."
378         else
379             echo " [!] tailscale up didn't respond (tailscaled may be wedged)"
380         fi
381
382         local ts_args=""
383         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
384         if run_with_timeout 10 "$ts_bin" down $ts_args >> output.log 2>&1; then
385             echo " [] Tailscale disconnected."
386         else
387             echo " [!] tailscale down didn't respond"
388         fi
389     else
390         echo " [!] tailscale CLI not found skipping"
391     fi
392 }
393
394 stop_pidfile_daemon() {
395     local pid_file="$1"
396     local label="$2"
397     if [ -f "$pid_file" ]; then
398         local pid
399         pid=$(cat "$pid_file")
400         if kill -0 "$pid" 2>/dev/null; then
401             echo " Stopping $label (PID $pid)..."
402             sudo -n kill "$pid" 2>> output.log || kill "$pid" 2>> output.log || true
403             sleep 0.5

```

```

403     if kill -0 "$pid" 2>/dev/null; then
404         sudo -n kill -9 "$pid" 2>> output.log || kill -9 "$pid" 2>> output.log || true
405     fi
406 fi
407 rm -f "$pid_file"
408 fi
409 }
410
411 disable_all_proxies() {
412     local svc="$1"
413     if [ -n "$svc" ]; then
414         sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
415         sudo -n networksetup -setwebproxystate "$svc" off 2>> output.log || true
416         sudo -n networksetup -setsecurewebproxystate "$svc" off 2>> output.log || true
417     fi
418 }
419
420 flush_dns_cache() {
421     if [[ "$OSTYPE" == "darwin"* ]]; then
422         sudo -n dscacheutil -flushcache 2>> output.log || true
423         sudo -n killall -HUP mDNSResponder 2>> output.log || true
424     else
425         resolvectl flush-caches 2>> output.log || true
426     fi
427 }
428
429 wait_for_dhcp_settle() {
430     # Resolve physical interface (Wi-Fi or Ethernet)
431     local physical_iface
432     physical_iface=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|
433     Ethernet)/{getline; print $2; exit}')
434     [ -z "$physical_iface" ] && physical_iface="en0"
435
436     echo -n "Waiting for physical network interface ($physical_iface) DHCP lease to settle..."
437     local settled=false
438     for attempt in {1..60}; do
439         PHYSICAL_GW=$(ipconfig getpacket "$physical_iface" 2>> output.log | awk -F'{}' '{print $2}'
440         )
441         local local_ip
442         local_ip=$(ifconfig "$physical_iface" 2>/dev/null | awk '/inet /{print $2}')
443
444         if [ -n "$PHYSICAL_GW" ] && [[ "$PHYSICAL_GW" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
445             [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
446             [[ "$local_ip" != "169.254.*" ]]; then
447             echo "settled (Interface IP: $local_ip, Router IP: $PHYSICAL_GW)."
448             settled=true
449             break
450         fi
451
452         if [ "$attempt" -gt 15 ] && [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$
453         ]] && [[ "$local_ip" != "169.254.*" ]]; then
454             local fallback_gw
455             fallback_gw=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
456             if [ -n "$fallback_gw" ] && [[ "$fallback_gw" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]]; then
457                 PHYSICAL_GW="$fallback_gw"
458                 fi
459                 echo "static/settled detected (Interface IP: $local_ip, Gateway IP: ${PHYSICAL_GW:-unknown})."
460                 settled=true
461                 break
462             fi
463             echo -n "."
464             sleep 0.5
465         done
466         if [ "$settled" = "false" ]; then
467             echo "timed out or no active link. Proceeding anyway."
468         fi
469     }
470
471 reset_sharing_services() {
472     if [[ "$OSTYPE" == "darwin"* ]]; then
473         sudo -n killall sharingd rapportd 2>> output.log || true
474     fi
475 }
476
477 read_adguard_ip() {
478     if [ -f ".gateway_ip" ]; then
479         tr -d '\r' < ".gateway_ip" | awk 'NR==1{print $1;exit}'
480     fi
481 }
482
483 is_kill_switch_enabled() {

```

```

481 [ "$KILL_SWITCH" = "true" ]
482 }
483
484 add_warp_bypass_routes() {
485     local target="${1:-}"
486     local ep1
487     ep1=$(get_warp_endpoint_1)
488     local ep2
489     ep2=$(get_warp_endpoint_2)
490
491     if [[ "$OSTYPE" == "darwin"* ]]; then
492         local routing_arg=""
493         local msg_via=""
494         if [ -n "$target" ]; then
495             if [[ "$target" =~ ^en[0-9]+$ || "$target" == "bridge"* ]]; then
496                 routing_arg="-interface $target"
497                 msg_via="interface $target"
498             else
499                 routing_arg="$target"
500                 msg_via="gateway $target"
501             fi
502         else
503             local gateway_ip
504             gateway_ip=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
505             if [ -z "$gateway_ip" ]; then
506                 local iface
507                 iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
508                 if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
509                     iface="en0"
510                 fi
511                 routing_arg="-interface $iface"
512                 msg_via="interface $iface"
513             else
514                 routing_arg="$gateway_ip"
515                 msg_via="gateway $gateway_ip"
516             fi
517         fi
518
519         echo -e "\nAdding host bypass routes for Cloudflare WARP endpoints via $msg_via..."
520         sudo -n route delete -host "$ep1" 2>/dev/null || true
521         sudo -n route delete -host "$ep2" 2>/dev/null || true
522         sudo -n route add -host "$ep1" $routing_arg >> output.log 2>&1 || true
523         sudo -n route add -host "$ep2" $routing_arg >> output.log 2>&1 || true
524
525         echo -e "Adding host bypass routes for Tailscale control plane via $msg_via..."
526         sudo -n route delete -net 192.200.0.0/24 2>/dev/null || true
527         sudo -n route add -net 192.200.0.0/24 $routing_arg >> output.log 2>&1 || true
528
529         echo -e "Adding host bypass routes for Tailscale DERP relays via $msg_via..."
530         local derp_ips
531         derp_ips=$(curl -s --connect-timeout 5 https://login.tailscale.com/derpmap/default | grep -oE 'IPv4
532         "[[:space:]]*:[[:space:]]*" [0-9.]+' | cut -d'"' -f4 | grep -E '^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' ||
533         true)
534         if [ -n "$derp_ips" ]; then
535             echo "$derp_ips" > /tmp/nullexit-derp-ips.txt
536             for ip in $derp_ips; do
537                 sudo -n route delete -host "$ip" 2>/dev/null || true
538                 sudo -n route add -host "$ip" $routing_arg >> output.log 2>&1 || true
539             done
540         fi
541     else
542         local gateway_ip="${target}"
543         if [ -z "$gateway_ip" ]; then
544             gateway_ip=$(ip route show default 2>/dev/null | awk '/default/ {print $3}')
545         fi
546         if [ -n "$gateway_ip" ]; then
547             sudo ip route add "$ep1" via "$gateway_ip" >> output.log 2>&1 || true
548             sudo ip route add "$ep2" via "$gateway_ip" >> output.log 2>&1 || true
549         fi
550     fi
551 }
552
553 remove_warp_bypass_routes() {
554     local ep1
555     ep1=$(get_warp_endpoint_1)
556     local ep2
557     ep2=$(get_warp_endpoint_2)
558
559     if [[ "$OSTYPE" == "darwin"* ]]; then
560         echo -e "\nRemoving host bypass routes for Cloudflare WARP endpoints..."
561         sudo -n route delete -host "$ep1" >> output.log 2>&1 || true

```

```

560 sudo -n route delete -host "$sep2" >> output.log 2>&1 || true
561 echo -e "Removing host bypass routes for Tailscale control plane..."
562 sudo -n route delete -net 192.200.0.0/24 >> output.log 2>&1 || true
563 if [ -f /tmp/nullexit-derp-ips.txt ]; then
564     echo -e "Removing host bypass routes for Tailscale DERP relays..."
565     while read -r ip; do
566         sudo -n route delete -host "$ip" >> output.log 2>&1 || true
567     done < /tmp/nullexit-derp-ips.txt
568     rm -f /tmp/nullexit-derp-ips.txt
569 fi
570 else
571     sudo ip route del "$sep1" >> output.log 2>&1 || true
572     sudo ip route del "$sep2" >> output.log 2>&1 || true
573 fi
574 }
575
576 # FaceTime / real-time P2P split-route (opt-in, default OFF)
577 # Mitigation for the FaceTime Double-Tunnel bug (15.12.22). Adds more-specific
578 # host routes for Apple's FaceTime media/relay /16s via the PHYSICAL gateway, so
579 # longest-prefix match beats the exit-node 0.0.0.0/1 + 128.0.0.0/1 and ONLY those
580 # /16s leave direct real-time media can then hole-punch from the real IP instead
581 # of dying behind WARP's symmetric NAT. Mirrors add_warp_bypass_routes.
582 #
583 # PRIVACY: the direct /16s expose the real ISP IP (not WARP) for that traffic
584 # a deliberate, narrow hole in the Double-Tunnel promise, scoped to Apple infra
585 # (which already ties traffic to your Apple ID). Inert unless FACETIME_SPLIT_ROUTE
586 # =true. Needs BOTH a route AND a PF pass: 'block out on en0 all' would otherwise
587 # drop the packets, so we populate the <apple_direct> table pf.conf permits.
588 facetime_split_enabled() {
589     [ "$(read_env_var FACETIME_SPLIT_ROUTE | tr '[:upper:]' '[:lower:]')" = "true" ]
590 }
591
592 facetime_direct_subnets() {
593     local s; s=$(read_env_var FACETIME_DIRECT_SUBNETS)
594     echo "${s:-17.249.0.0/16}"
595 }
596
597 add_apple_split_routes() {
598     facetime_split_enabled || return 0
599     [[ "$OSTYPE" == "darwin"* ]] || return 0
600     local gw="${1:-}"
601     if [ -z "$gw" ]; then
602         # The physical default route survives alongside the exit-node /1 override,
603         # so the real gateway is still the 'default' entry in the table.
604         gw=$(netstat -rn -f inet 2>/dev/null | awk '1=="default"{print $2; exit}')
605     fi
606     if [ -z "$gw" ] || [[ "$gw" =~ ^utun ]] || [[ "$gw" =~ ^link ]]; then
607         echo " [FaceTime split-route] could not resolve a physical gateway skipping."
608         return 0
609     fi
610     local subnets; subnets=$(facetime_direct_subnets)
611     echo -e "\n[FaceTime split-route] routing Apple subnets DIRECT via $gw (real IP exposed for these):
612     $subnets"
613     for net in $subnets; do
614         sudo -n route delete -net "$net" >> output.log 2>&1 || true
615         sudo -n route add -net "$net" "$gw" >> output.log 2>&1 || true
616         # Kill-switch pass: the <apple_direct> table is what pf.conf lets out on en0.
617         sudo -n pfctl -a com.apple/nullexit -t apple_direct -T add "$net" >> output.log 2>&1 || true
618     done
619 }
620
621 remove_apple_split_routes() {
622     [[ "$OSTYPE" == "darwin"* ]] || return 0
623     local subnets; subnets=$(facetime_direct_subnets)
624     for net in $subnets; do
625         sudo -n route delete -net "$net" >> output.log 2>&1 || true
626     done
627     # Flush the whole table so no direct Apple exception lingers after teardown.
628     sudo -n pfctl -a com.apple/nullexit -t apple_direct -T flush >> output.log 2>&1 || true
629 }
630
631 bounce_wifi_interfaces() {
632     if [[ "$OSTYPE" == "darwin"* ]]; then
633         local wifi_port
634         wifi_port=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline;
635         print $2}')
636         if [ -n "$wifi_port" ]; then
637             sudo -n networksetup -setairportpower "$wifi_port" off 2>> output.log || true
638             sleep 2
639             sudo -n networksetup -setairportpower "$wifi_port" on 2>> output.log || true
640             echo "Wi-Fi bounced (interface $wifi_port)"
641         fi
642     fi
643 }

```

```

639     sleep 3
640 else
641     echo "Could not detect Wi-Fi interface. Bouncing fallback interfaces..."
642     set +e
643     for iface in en0 en1 en2 en3 en4 en5; do
644         if ifconfig "$iface" >> output.log 2>&1; then
645             sudo -n ifconfig "$iface" down 2>> output.log || true
646             sudo -n ifconfig "$iface" up 2>> output.log || true
647         fi
648     done
649     set -e
650     echo "Fallback interfaces bounced"
651 fi
652 fi
653 }
654
655 get_active_interface() {
656     local iface
657     iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
658
659     # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical
660     interface
661     if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
662         for i in en0 en1 en2 en3; do
663             if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -
664                 q "inet "; then
665                 iface="$i"
666                 break
667             fi
668         done
669     fi
670
671     # Default fallback if still empty or not enX
672     if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
673         iface="en0"
674     fi
675
676     echo "$iface"
677 }
678
679 enable_killswitch() {
680     if [[ "$OSTYPE" == "darwin"* ]]; then
681         if ! is_kill_switch_enabled; then
682             return 0
683         fi
684
685         local ext_if
686         ext_if=$(get_active_interface)
687         local ep1
688         ep1=$(get_warp_endpoint_1)
689         local ep2
690         ep2=$(get_warp_endpoint_2)
691         local pf_conf_path
692         pf_conf_path="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/pf.conf"
693
694         # Keep the host's scrub max-mss in sync with GATEWAY_MSS (same variable that drives
695         # post-rules.txt's --set-mss for the container). The host egresses through the same
696         # double-tunnel (Tailscale+WARP) as forwarded phone traffic, so it needs the same
697         # empirically-safe clamp a stale/higher value here reproduces the 15.3.2 stalling
698         # bug for the host's own connections (see devref.md).
699         local gateway_mss
700         gateway_mss=$(read_env_var GATEWAY_MSS)
701         if [[ "$gateway_mss" =~ ^[0-9]+$ ]]; then
702             sed -i '' "s/max-mss [0-9]*/max-mss ${gateway_mss}/g" "$pf_conf_path" 2>/dev/null || true
703         fi
704
705         echo -e "\nEnabling macOS Packet Filter (PF) Kill-Switch on $ext_if..."
706
707         # Enable PF and capture the reference token
708         local token_out
709         token_out=$(sudo -n pfctl -E 2>> output.log || true)
710         local token
711         token=$(echo "$token_out" | awk '/Token/{print $NF}')
712         if [ -n "$token" ]; then
713             echo "$token" > /tmp/nullexit-pf-token.txt
714             echo " [] PF enabled with token $token."
715         else
716             echo " [!] WARNING: Could not capture PF reference token (may already be enabled or sudo failed)."
717         fi
718
719         # Load the rules into the anchor

```



```

718 if sudo -n pfctl -D "ext_if=$ext_if" -a com.apple/nullexit -f "$pf_conf_path" >> output.log 2>&1;
719 then
720     echo " [] Ruleset loaded into anchor com.apple/nullexit."
721 else
722     echo " [!] ERROR: Failed to load ruleset into anchor."
723 fi
724
725 # Populate tables in the anchor
726 sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T flush >> output.log 2>&1 || true
727 sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep1" >> output.log 2>&1 || true
728 sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep2" >> output.log 2>&1 || true
729
730 if [ -f /tmp/nullexit-derp-ips.txt ]; then
731     local derp_ips
732     derp_ips=$(tr '\n' ' ' < /tmp/nullexit-derp-ips.txt)
733     if [ -n "$derp_ips" ]; then
734         sudo -n pfctl -a com.apple/nullexit -t derp_relays -T flush >> output.log 2>&1 || true
735         sudo -n pfctl -a com.apple/nullexit -t derp_relays -T add $derp_ips >> output.log 2>&1 || true
736     fi
737 fi
738 echo " [] Kill-switch tables populated."
739 fi
740
741 disable_killswitch() {
742     if [[ "$OSTYPE" == "darwin"* ]]; then
743         if ! is_kill_switch_enabled; then
744             return 0
745         fi
746         echo -e "\nDisabling macOS PF Kill-Switch..."
747
748         # Flush rules from anchor com.apple/nullexit
749         sudo -n pfctl -a com.apple/nullexit -F all >> output.log 2>&1 || true
750
751         # Release the enable reference if token is present
752         if [ -f /tmp/nullexit-pf-token.txt ]; then
753             local token
754             token=$(cat /tmp/nullexit-pf-token.txt)
755             if [ -n "$token" ]; then
756                 sudo -n pfctl -X "$token" >> output.log 2>&1 || true
757                 echo " [] Released PF enable reference token $token."
758             fi
759             rm -f /tmp/nullexit-pf-token.txt
760         else
761             echo " [] Rules flushed from anchor com.apple/nullexit."
762         fi
763     fi
764 }
765
766 write_host_ips() {
767     # Gather all active IPv4 addresses on the host and write to .host_ips
768     # routing-fix.sh will read this file to explicitly whitelist them for direct Tailscale P2P
769     local repo_dir
770     repo_dir="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"
771
772     if [[ "$OSTYPE" == "darwin"* ]]; then
773         ifconfig | awk '/inet /{print $2}' | grep -v '127.0.0.1' > "$repo_dir/.host_ips" 2>/dev/null || true
774     else
775         ip -4 addr show 2>/dev/null | awk '/inet /{print $2}' | cut -d/ -f1 | grep -v '127.0.0.1' > "
776         $repo_dir/.host_ips" 2>/dev/null || true
777     fi
778 }
779
780 start_dns_watcher() {
781     if [[ "$OSTYPE" == "darwin"* ]]; then
782         local target_ip=$1
783         stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
784         echo " Starting background DNS Watcher for seamless Wi-Fi roaming..."
785         nohup bash -c "
786             source \"\$SCRIPT_DIR/scripts/common.sh\"
787             trap 'exit 0' SIGTERM SIGINT SIGHUP
788             while true; do
789                 ACTIVE_IF=$(get_active_service)
790                 if [ -n \"\$ACTIVE_IF\" ]; then
791                     CURRENT_DNS=$(networksetup -getdnsservers \"\$ACTIVE_IF\" 2>/dev/null)
792                     if [ \"\$CURRENT_DNS\" != \"\$target_ip\" ]; then
793                         networksetup -setdnsservers \"\$ACTIVE_IF\" \"\$target_ip\" >/dev/null 2>&1
794                     fi
795                 fi
796                 sleep 30

```

```

797     done
798     " >> output.log 2>&1 &
799     echo $! > "$DNS_WATCHER_PID_FILE"
800 else
801     local target_ip=$1
802     stop_dns_watcher
803     echo " Starting background DNS Watcher for seamless Wi-Fi roaming..."
804     nohup bash -c "
805         trap 'exit 0' SIGTERM SIGINT SIGHUP
806         while true; do
807             if [ -n \"\$ACTIVE_SERVICE\" ]; then
808                 CURRENT_DNS=\$(resolvectl dns \"\$ACTIVE_SERVICE\" 2>/dev/null | grep -oE
809                 '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head -1)
810                 if [ \"\$CURRENT_DNS\" != \"\$target_ip\" ]; then
811                     resolvectl dns \"\$ACTIVE_SERVICE\" \"\$target_ip\" >/dev/null 2>&1 || true
812                 fi
813             fi
814             sleep 30
815         done
816     " >> output.log 2>&1 &
817     echo $! > "$DNS_WATCHER_PID_FILE"
818 fi
819 }
820 # SOCKS5 Proxy (WARP Tunnel)
821 # When Tailscale's exit node can't route through WARP (due to userspace
822 # forwarding), we use a SOCKS5 proxy running inside the warp container's
823 # network namespace. Connections created by this proxy go through the
824 # kernel routing table (table 200 -> tun0 -> WARP), unlike Tailscale's
825 # userspace exit node which bypasses it.
826 #
827 # The SOCKS5 proxy is exposed on localhost:${SOCKS_PROXY_PORT} via Docker port mapping.
828 # 'networksetup -setsocksfirewallproxy' on macOS routes system TCP traffic
829 # through the proxy, which then goes through WARP.
830 #
831 # IMPORTANT: CLI tools like curl do NOT respect macOS system proxy settings.
832 # They must use --socks5-hostname or the ALL_PROXY env variable explicitly.
833
834 enable_socks_proxy() {
835     if [[ "$OSTYPE" == "darwin"* ]]; then
836         local svc="$1"
837         if [ -z "$svc" ]; then
838             svc="$ACTIVE_SERVICE"
839         fi
840
841         echo -n " Enabling system-wide SOCKS5 proxy on $svc (localhost:${SOCKS_PROXY_PORT})... "
842         sudo -n networksetup -setsocksfirewallproxy "$svc" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.log || true
843         sudo -n networksetup -setsocksfirewallproxystate "$svc" on 2>> output.log || true
844
845         # Also set on en0 service for macOS resolver consistency
846         if [ "$svc" != "$ENO_SERVICE" ]; then
847             sudo -n networksetup -setsocksfirewallproxy "$ENO_SERVICE" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.
848             log || true
849             sudo -n networksetup -setsocksfirewallproxystate "$ENO_SERVICE" on 2>> output.log || true
850         fi
851
852         # IMPORTANT: Exclude local LAN and mDNS traffic from the proxy so P2P works natively
853         # without source-IP masking from the Colima VM bridge.
854         sudo -n networksetup -setproxybypassdomains "$svc" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
855         172.16.0.0/12 "*.local" 2>> output.log || true
856         if [ "$svc" != "$ENO_SERVICE" ]; then
857             sudo -n networksetup -setproxybypassdomains "$ENO_SERVICE" 127.0.0.1 localhost 192.168.0.0/16
858             10.0.0.0/8 172.16.0.0/12 "*.local" 2>> output.log || true
859         fi
860
861         echo "done."
862         echo " All TCP traffic now routed through gateway -> WARP tunnel -> internet."
863         echo " (CLI tools: export ALL_PROXY=socks5://127.0.0.1:$SOCKS_PROXY_PORT)"
864     else
865         local svc="$1"
866         echo " (Linux global SOCKS proxy configuration is desktop-environment specific, skipping.)"
867         echo " SOCKS5 proxy is active and listening on 127.0.0.1:$SOCKS_PROXY_PORT"
868     fi
869 }
870
871 disable_socks_proxy() {
872     if [[ "$OSTYPE" == "darwin"* ]]; then
873         for svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
874             sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
875         done
876         echo " SOCKS5 proxy disabled."
877     fi
878 }

```

```

874 else
875     local svc="$1"
876     echo " (Linux global SOCKS proxy configuration skipped.)"
877 fi
878 }
879
880 # Helper to restore host DNS to a clean state (used on script start, on failure,
881 # and after a successful STOP). Without this, a successful toggle-off leaves the
882 # host's resolver pinned to a now-dead gateway IP every lookup stalls for macOS's
883 # DNS timeout (~5s) before falling through to whatever next server is configured.
884 # With it, we always leave the host in '1.1.1.1 / no search domain'.
885 #
886 # We set DNS on BOTH the active service AND the en0 service because macOS's scutil DNS
887 # resolver is commonly scoped to en0 (Wi-Fi). A per-service change on e.g.
888 # "USB 10/100 LAN" is ignored by the system resolver unless the en0 service is also updated.
889 reset_dns() {
890     if [[ "$OSTYPE" == "darwin"* ]]; then
891         if [ -n "$ACTIVE_SERVICE" ]; then
892             for dns_svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
893                 networksetup -setsearchdomains "$dns_svc" "Empty" 2>> output.log || true
894                 networksetup -setdnsservers "$dns_svc" 1.1.1.1 2>> output.log || true
895             done
896         fi
897     else
898         if [ -n "$ACTIVE_SERVICE" ]; then
899             resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
900             resolvectl revert "$ACTIVE_SERVICE" 2>/dev/null || true
901         fi
902     fi
903 }
904
905 # Verify basic direct-internet connectivity after a recovery/teardown: DNS via
906 # 1.1.1.1, then HTTP (curl) or ping to 1.1.1.1. Sets the global INTERNET_OK to
907 # "true" on any successful check (callers read $INTERNET_OK in their summary).
908 # Always returns 0 so a bare call can't trip 'set -e'. Used by recover.sh's macOS
909 # and Linux recovery paths (extracted from a formerly-duplicated block).
910 verify_internet_connectivity() {
911     step "Verifying internet connectivity"
912
913     # Wait a moment for the network to settle
914     sleep 2
915
916     INTERNET_OK=false
917
918     # Try DNS resolution first
919     if host -W 3 google.com 1.1.1.1 >> output.log 2>&1; then
920         ok "DNS resolution works (google.com via 1.1.1.1)"
921         INTERNET_OK=true
922     elif nslookup google.com 1.1.1.1 >> output.log 2>&1; then
923         ok "DNS resolution works (nslookup google.com via 1.1.1.1)"
924         INTERNET_OK=true
925     else
926         warn "DNS resolution check failed will still try ping/curl"
927     fi
928
929     # Try actual HTTP connectivity
930     if command -v curl >> output.log 2>&1; then
931         if curl -sf --max-time 5 https://1.1.1.1 >> output.log 2>&1; then
932             ok "Internet reachable via HTTP"
933             INTERNET_OK=true
934         elif curl -sf --max-time 5 http://1.1.1.1 >> output.log 2>&1; then
935             ok "Internet reachable via HTTP (plain)"
936             INTERNET_OK=true
937         else
938             warn "HTTP check failed"
939         fi
940     elif command -v ping >> output.log 2>&1; then
941         if ping -c 1 -W 3 1.1.1.1 >> output.log 2>&1; then
942             ok "Internet reachable via ping"
943             INTERNET_OK=true
944         else
945             warn "Ping check failed"
946         fi
947     fi
948
949     return 0
950 }
951
952 # Local DNS Proxy (Python)
953 # When the tailnet data plane cannot establish, we use a Python DNS proxy.
954 # It listens on UDP:53, receives DNS queries, forwards them over TCP to

```

```

955 # AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (Docker port mapping), and returns the response.
956 # The Python script handles the DNS-over-TCP wire format (2-byte length prefix)
957 # correctly unlike socat which just forwards raw bytes.
958 #
959 # Python3 is built into macOS no external dependencies.
960 DNS_PROXY_PID=""
961
962 # Kill any leftover DNS proxy process
963 stop_local_dns_proxy() {
964     if [ -n "$DNS_PROXY_PID" ]; then
965         kill "$DNS_PROXY_PID" 2>/dev/null || true
966         wait "$DNS_PROXY_PID" 2>/dev/null || true
967         DNS_PROXY_PID=""
968     fi
969     # Clean up any stale Python DNS proxy processes
970     sudo -n pkill -f "dns-proxy.py" 2>/dev/null || true
971 }
972
973 # Start local DNS proxy (Python)
974 start_local_dns_proxy() {
975     if [ -z "$DNS_PROXY_BIN" ]; then
976         echo " Python not found. Python3 is built into macOS this shouldn't happen."
977         return 1
978     fi
979
980     if [ ! -f "$DNS_PROXY_SCRIPT" ]; then
981         echo " DNS proxy script not found at $DNS_PROXY_SCRIPT"
982         return 1
983     fi
984
985     # Kill any leftover process first
986     stop_local_dns_proxy
987
988     echo -n " Starting local DNS proxy via Python (UDP:53 TCP:${DNS_PROXY_PORT})... "
989     sudo -n bash -c "TARGET_PORT=$DNS_PROXY_PORT \"$DNS_PROXY_BIN\" \"$DNS_PROXY_SCRIPT\" &
990     DNS_PROXY_PID=$!
991     disown "$DNS_PROXY_PID" 2>/dev/null || true
992
993     # Give it a moment to bind
994     sleep 0.5
995
996     if kill -0 "$DNS_PROXY_PID" 2>/dev/null; then
997         echo "started (PID $DNS_PROXY_PID).\"
998
999         # Hijack host DNS to localhost
1000         echo -n " Hijacking host DNS to 127.0.0.1 for ad-blocking... \"
1001         if [[ \"$OSTYPE\" == \"darwin*\" ]]; then
1002             networksetup -setsearchdomains \"$ACTIVE_SERVICE\" \"Empty\" 2>/dev/null || true
1003             networksetup -setsearchdomains \"$ENO_SERVICE\" \"Empty\" 2>/dev/null || true
1004             if ! networksetup -setdnsservers \"$ACTIVE_SERVICE\" 127.0.0.1; then
1005                 echo \"FAILED (networksetup error check permissions).\"
1006                 stop_local_dns_proxy
1007                 return 1
1008             fi
1009             networksetup -setdnsservers \"$ENO_SERVICE\" 127.0.0.1 2>/dev/null || true
1010             sudo -n dscacheutil -flushcache 2>/dev/null || true
1011             sudo -n killall -HUP mDNSResponder 2>/dev/null || true
1012         else
1013             resolvectl domain \"$ACTIVE_SERVICE\" \"\" 2>/dev/null || true
1014             resolvectl domain \"$ACTIVE_SERVICE\" \"\" 2>/dev/null || true
1015             if ! sudo -n resolvectl dns \"$ACTIVE_SERVICE\" 127.0.0.1; then
1016                 echo \"FAILED (resolvectl error check permissions).\"
1017                 stop_local_dns_proxy
1018                 return 1
1019             fi
1020             resolvectl dns \"$ACTIVE_SERVICE\" 127.0.0.1 2>/dev/null || true
1021             sudo -n resolvectl flush-caches 2>/dev/null || true
1022         fi
1023
1024         # Verify DNS actually works through the proxy
1025         echo -n \" Verifying DNS resolution... \"
1026         if dig @127.0.0.1 google.com +short +timeout=5 &>/dev/null; then
1027             echo \"ok.\"
1028             return 0
1029         else
1030             echo \"FAILED (DNS queries not reaching AdGuard).\"
1031             echo \" Restoring DNS to 1.1.1.1...\"
1032             reset_dns
1033             stop_local_dns_proxy
1034             return 1
1035         fi

```

```

1036 else
1037     echo "FAILED (could not bind port 53 is it in use?)."
1038     DNS_PROXY_PID=""
1039     return 1
1040 fi
1041 }
1042
1043 # Helper to FORCIBLY hijack host DNS to a single server (the gateway's AdGuard IP)
1044 # and VERIFY via 'networksetup -getdnsservers' that the only name server is exactly
1045 # $1. Returns 0 iff the live state matches; non-zero if it can't be confirmed.
1046 #
1047 # This deliberately does NOT append a 1.1.1.1 fallback: macOS's resolver queries
1048 # the list in order and falls back to the next entry on timeout, so on a
1049 # misbehaving AdGuard (filter compile crash, OOM kill, etc.) the resolver still
1050 # leaks to 1.1.1.1 and silently bypasses ad-blocking. With no fallback, a broken
1051 # gateway manifests as a *visible* DNS outage that prompts the user to
1052 # investigate which is the correct failure mode.
1053 force_dns_to_gateway() {
1054     if [[ "$OSTYPE" == "darwin*" ]]; then
1055         local target_ip="$1"
1056         if [ -z "$target_ip" ] || [ -z "$ACTIVE_SERVICE" ]; then
1057             echo " [force_dns] ERROR: target_ip or ACTIVE_SERVICE is empty" >&2
1058             return 1
1059         fi
1060
1061         local attempt
1062         for attempt in 1 2 3; do
1063             echo -n " [force_dns] Attempt $attempt/3: setting DNS to $target_ip on \"\$ACTIVE_SERVICE\" + \"\$ENO_SERVICE\"... "
1064
1065             # Apply to BOTH the active service AND the en0 service the scutil DNS resolver
1066             # is scoped to en0 (Wi-Fi) and ignores per-service changes that don't include it.
1067             # Fail fast on ACTIVE_SERVICE; best-effort on ENO_SERVICE (it may not exist).
1068             networksetup -setsearchdomains "$ACTIVE_SERVICE" "ts.net" || true
1069             if ! networksetup -setdnsservers "$ACTIVE_SERVICE" "$target_ip"; then
1070                 echo "FAILED (networksetup error check permissions)"
1071                 sleep 1
1072                 continue
1073             fi
1074             networksetup -setsearchdomains "$ENO_SERVICE" "ts.net" || true
1075             networksetup -setdnsservers "$ENO_SERVICE" "$target_ip" || true
1076
1077             local entries
1078             entries=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log \
1079                 | awk '/^(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)$/ {print}')
1080             if echo "$entries" | grep -qFx "$target_ip"; then
1081                 echo "VERIFIED"
1082                 return 0
1083             fi
1084
1085             local current
1086             current=$(echo "$entries" | tr '\n' ' ' | sed 's/ $//')
1087             echo "NOT YET (current DNS: [${current:-none}])"
1088
1089             if [ "$attempt" = "3" ]; then sleep 2; else sleep 1; fi
1090         done
1091
1092         echo " [force_dns] FAILED after 3 attempts"
1093         return 1
1094     else
1095         local target_ip="$1"
1096
1097         for attempt in {1..3}; do
1098             echo -n " [force_dns] Attempt $attempt/3: setting DNS to $target_ip on $ACTIVE_SERVICE... "
1099             sudo -n resolvectl domain "$ACTIVE_SERVICE" "~ts.net" 2>/dev/null || true
1100             if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" "$target_ip"; then
1101                 echo "FAILED (resolvectl error)"
1102             else
1103                 # Verify
1104                 entries=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' )
1105                 if [ "$entries" == "$target_ip" ]; then
1106                     echo "VERIFIED"
1107                     return 0
1108                 else
1109                     echo "MISMATCH (Resolver refused lock)"
1110                 fi
1111             fi
1112             sleep 2
1113         done

```

```

1114     return 1
1115 fi
1116 }

```

12 scripts/watcher.sh

```

1 #!/bin/bash
2 # scripts/watcher.sh nullexit post-wake / post-roam recovery watcher
3 #
4 # Long-running daemon. Launched by launchd as a LaunchAgent at
5 # ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
6 # (label com.nullexit.wake-recovery).
7 #
8 # Two independent listeners run as backgrounded pipelines:
9 #
10 # (1) WAKE listener 'log stream' live-follows the macOS unified log for
11 # power-management events (lid closewake, manual wake, DarkWake from
12 # clamshell). Each event triggers run_recover (with a global debounce so
13 # a single wake that emits 2-3 log lines only causes ONE post-wake run).
14 #
15 # (2) NETWORK listener 'scutil n.watch' on 'State:/Network/Global/IPv4'
16 # fires on every Wi-Fi roam, ethernet swap, hotspot join, captive-portal
17 # re-bind, VPN up/down, etc. Same debounce.
18 #
19 # Both listeners call run_recover, which:
20 # - checks /tmp/nullexit-gateway-active.marker (only act when toggle.sh left
21 # the gateway up)
22 # - debounces (10s default) so multiple events don't pile up
23 # - shells out to 'bash <repo>/recover.sh --post-wake' directly (no GUI auth prompt needed due to
24 # NOPASSWD sudoers)
25 #
26 # See devref.md 10.29 for the full Apple power management primer and the
27 # rationale behind using 'log stream' + 'scutil n.watch' from a shell daemon.
28 #
29 # NOTE: errexit ('set -e') is intentionally NOT enabled. This is a long-running
30 # daemon, not a discrete-step script, and errexit actively breaks the reconnect
31 # loops below: any failed pipe (log stream on rotation, scutil on state churn)
32 # would kill the outer subshell before 'echo "reconnecting..."; sleep 5;' runs,
33 # AND the parent's final 'wait' would propagate a non-zero child exit and kill
34 # the entire watcher. Every error path is already wrapped in '|| true' /
35 # '2>/dev/null' fallbacks.
36 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
37 source "$SCRIPT_DIR/common.sh"
38 setup_standard_path
39
40 LOG="$SCRIPT_DIR/./output.log"
41 # Resolve our own directory; recover.sh lives one level up at the repo root.
42 # This makes the daemon install-agnostic no need to template the path
43 # in launchd, no need to keep a deploy-specific hardcoded location in sync.
44 RECOVER="$SCRIPT_DIR/./recover.sh"
45 MARKER="$MARKER_FILE"
46 DEBOUNCE_FILE="/tmp/nullexit-watcher.last-recovery"
47 DEBOUNCE_SECONDS="${NULLEXIT_DEBOUNCE_SECONDS:-10}"
48 # Post-recovery SETTLE cooldown. A post-wake run re-asserts the exit node
49 # (recover.sh: 'tailscale set --exit-node'), which re-writes the host default
50 # route and that route IS State:/Network/Global/IPv4, the exact key Listener 2
51 # watches. The self-triggered change lands ~15-45s AFTER the run finishes, past
52 # the 10s debounce, so without a longer window the watcher re-fires post-wake in
53 # a ~40s self-feedback loop that only stops once routing happens to reach a fixed
54 # point (devref 15.12.21). This cooldown is armed ONLY after a recovery runs, so
55 # a genuine first wake/roam is still handled immediately; only the change we
56 # caused ourselves gets absorbed. A real new roam during the window is delayed at
57 # most this many seconds, which is acceptable.
58 COOLDOWN_FILE="/tmp/nullexit-watcher.settle-until"
59 POSTWAKE_SETTLE_SECONDS="${NULLEXIT_POSTWAKE_SETTLE_SECONDS:-45}"
60
61 exec >> "$LOG" 2>&1
62
63 # Single-instance lock
64 # Without this, repeated lid-closewake cycles under macOS Sonoma+ accumulate
65 # orphan watcher.sh processes: launchd suspends on sleep and SIGCONT-resumes
66 # on wake, and KeepAlive.Crashed=true caused relaunch on any non-zero exit;
67 # the listener grandchildren (log stream | while ... and (echo n.add ...;
68 # sleep 86400) | scutil | while ...) stay alive because pipe producers don't
69 # reliably respond to plain SIGTERM if the consumer is in a half-closed state.
70 #

```

```

70 # We solve it with a PID-file lock (POSIX-portable; macOS does NOT ship 'flock').
71 # The trap below removes the lock file on every exit so a stale lock never
72 # outlives a crashed watcher. On launch: read the existing PID, check it is
73 # alive AND is actually a 'scripts/watcher.sh' process; if so, exit 0; else
74 # overwrite with our PID and proceed. The check-then-write window has a TOCTOU
75 # race but is microseconds wide and only matters for hand-launched duplicate
76 # test invocations; production is single-launch via launchctl load -w.
77 LOCK_FILE="/tmp/nullexit-watcher.lock"
78 LOCK_PID=""
79 if [ -e "$LOCK_FILE" ]; then
80     LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
81     if [ -n "$LOCK_PID" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
82         # Verify the holder is actually a scripts/watcher.sh process (not some
83         # unrelated process that got the recycled PID). The exact-launchd pattern
84         # 'bash <abs-path>/scripts/watcher.sh' is hard to accidentally match with
85         # a 'grep scripts/watcher.sh .' invocation because we anchor on '^bash'
86         # followed by whitespace and end-of-line on the script path.
87         if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -qE "(^[[:space:]])bash[[:space:]]+.*scripts/
            watcher\\.sh$"; then
88             echo "[$(date -u +%FT%TZ)] watcher.sh: another instance holds $LOCK_FILE (pid=$LOCK_PID), exiting"
89             exit 0
90         fi
91     fi
92 fi
93 echo $$ > "$LOCK_FILE"
94
95 echo "[$(date -u +%FT%TZ)] watcher.sh start (pid=$$ ppid=$PPID user=$(id -un) lock=acquired)"
96
97 # Treat launchd-resume as a wake signal. Apple's launchd SUSPENDS LaunchAgents
98 # during system sleep and restores them on wake; during that suspended gap,
99 # macOS's 'log stream' cannot catch the wake event that prompted the resume.
100 # The next thing that happens after wake IS our own startup, so we always
101 # fire one post-wake run unconditionally here. The marker check + debounce
102 # in run_recover() make this idempotent: if the gateway isn't up or we just
103 # debounced, it no-ops. Without this, the FIRST wake-after-install goes
104 # unseen and the watcher appears broken until the SECOND wake.
105 # Helper: run recover.sh --post-wake if the gateway is active
106 run_recover() {
107     local why="$1"
108     if [ ! -f "$MARKER" ]; then
109         echo "[$(date -u +%FT%TZ)] $why no $MARKER, skip"
110         return 0
111     fi
112     local now last settle_until
113     now=$(date +%s)
114     # Post-recovery settle window: ignore the Global/IPv4 change our own last
115     # post-wake run caused (exit-node re-assert) so we don't self-feedback loop.
116     settle_until=$(cat "$COOLDOWN_FILE" 2>/dev/null || echo 0)
117     if [ "$now" -lt "$settle_until" ]; then
118         echo "[$(date -u +%FT%TZ)] $why settling ($(settle_until - now)s left in post-recovery cooldown),
            skip"
119         return 0
120     fi
121     last=$(cat "$DEBOUNCE_FILE" 2>/dev/null || echo 0)
122     if [ "$((now - last))" -lt "$DEBOUNCE_SECONDS" ]; then
123         echo "[$(date -u +%FT%TZ)] $why debounced ($(now - last)s since last, threshold ${DEBOUNCE_SECONDS}
            s)"
124         return 0
125     fi
126     echo "$now" > "$DEBOUNCE_FILE"
127     echo "[$(date -u +%FT%TZ)] $why bash $RECOVER --post-wake"
128     bash "$RECOVER" --post-wake
129     local exit_code=$?
130     now=$(date +%s)
131     echo "$now" > "$DEBOUNCE_FILE"
132     # Arm the settle cooldown so the route change this run just caused can't re-fire us.
133     echo "$((now + POSTWAKE_SETTLE_SECONDS))" > "$COOLDOWN_FILE"
134     echo "[$(date -u +%FT%TZ)] $why exit=$exit_code (settle until +${POSTWAKE_SETTLE_SECONDS}s, now=$now)"
135     return $exit_code
136 }
137
138 if [ -f "$MARKER" ]; then
139     run_recover "initial post-wake (launchd resume = wake signal)"
140 fi
141
142 # LAN P2P Auto-Detection
143 # Detects whether the current network allows safe Tailscale LAN P2P by:
144 # 1. Checking WPA2-Enterprise (802.1x) via system_profiler SPAirPortDataType always forces P2P=false
145 # because enterprise networks almost universally have AP Client Isolation.
146 # 2. AP Isolation probe sends a broadcast ping and checks if ANY LAN host responds.
147 # If no LAN hosts respond on a non-empty network, AP Isolation is active.

```



```

148 # Writes 'true' or 'false' to .lan_p2p_detected in the repo root.
149 # routing-fix.sh reads this file dynamically every 30s loop iteration.
150 P2P_OVERRIDE_FILE="$SCRIPT_DIR/../../lan_p2p_detected"
151
152
153 detect_lan_p2p_mode() {
154     local reason="unknown" allow="false"
155
156     # Step 1: Detect Wi-Fi security type via system_profiler (airport binary removed in macOS 14.4+)
157     # SPAirPortDataType lists all known networks; the first "Security:" line is the current association.
158     local security
159     security=$(system_profiler SPAirPortDataType 2>/dev/null \
160         | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF; exit}')
161
162     if [ -z "$security" ]; then
163         # Not on Wi-Fi or system_profiler unavailable fail safe
164         reason="no-wifi" allow="false"
165     elif echo "$security" | grep -qi "Enterprise"; then
166         # WPA2/WPA3 Enterprise = 802.1x always AP isolated
167         reason="802.1x-enterprise" allow="false"
168     elif echo "$security" | grep -qi "Personal\|None\|Open"; then
169         # WPA2/WPA3 Personal or open probe for AP isolation
170         local gw
171         gw=$(route -n get 1.1.1.1 2>/dev/null | awk '/gateway:/{print $2}')
172         if [ -z "$gw" ]; then
173             reason="no-default-route" allow="false"
174         else
175             local responses
176             responses=$(ping -c 3 -t 1 -q "$gw" 2>/dev/null | awk '/packets transmitted/{print $4}')
177             if [ "${responses:-0}" -gt 0 ]; then
178                 reason="wpa-personal-lan-reachable" allow="true"
179             else
180                 reason="wpa-personal-ap-isolated" allow="false"
181             fi
182         fi
183     else
184         # Unknown security type fail safe
185         reason="unknown-security-type" allow="false"
186     fi
187
188     echo "[$(date -u +%FT%TZ)] P2P detect: security='${security:-none}' allow=${allow} (reason: ${reason})"
189     echo "$allow" > "$P2P_OVERRIDE_FILE"
190     write_host_ips
191 }
192
193 # Run once on startup
194 detect_lan_p2p_mode
195
196 # Listener 1: WAKE events via unified log
197 # Predicate picks up:
198 # * "didWake" regular kernel wake events (lid open, RTC alarm)
199 # * "Wake from Sleep" powerd-driven normal wake
200 # * "DarkWake" clamshell-display wake that didn't fully wake the
201 # Mac (relevant when lid opens during display sleep)
202 # * "Waking from" generic catch-all for variations in newer macOS
203 # '[c]' makes the match case-insensitive (the unified log uses mixed-case
204 # phrases). --info reduces noise by dropping debug/trace log levels.
205 # Subsystem-based predicate catches every powermanagement emission (willSleep,
206 # didWake, didDim, didUndim, DarkWake, etc.) regardless of message phrasing
207 # across kernel versions. Phrasing-based predicates ('didWake', 'Wake from Sleep')
208 # are fragile across macOS releases. We accept the extra log noise because the
209 # run_recover() debounce + marker check filter out anything that isn't a real
210 # gateway-relevant event.
211 (
212     # Reconnect loop: macOS rotates the unified log buffer periodically; when
213     # that happens, this 'log stream' subscriber's pipe EOFs, the inner
214     # 'while read' consumer terminates, and without this wrapper the listener
215     # would be silently dead for the rest of the watcher's lifetime.
216     while ;; do
217         log stream --predicate \
218             ' subsystem == "com.apple.powermanagement" ' \
219             --style compact --info 2>/dev/null \
220             | while IFS= read -r line; do
221                 [ -z "$line" ] && continue
222                 run_recover "WAKE: $(echo "$line" | head -c 100)"
223             done
224         echo "[$(date -u +%FT%TZ)] log stream subshell exited; reconnecting in 5s"
225         sleep 5
226     done
227 ) &

```

```

228
229 # Listener 2: NETWORK state changes via scutil
230 # 'scutil n.watch' on a state key prints SCEventUpdate lines whenever the key
231 # changes. We watch Global/IPv4 because that's the umbrella that catches link,
232 # address, route, and DNS changes for ALL interfaces in one namespace.
233 # The pipe is kept alive by an 'sleep 86400' loop so scutil never sees EOF
234 # from us (which would silently terminate the watch).
235 (
236 # Reconnect loop: scutil can exit noisily on certain state changes (rare,
237 # but observed). Without the wrapper, the network listener's pipe EOFs and
238 # every subsequent roam goes unseen until launchd restarts us.
239 while ;; do
240 (
241 # Watch both IPv4 and IPv6 umbrella state IPv6-only or dual-stack
242 # networks never fire on the IPv4 key alone, and the user's first roam
243 # in a hotel / airport / on cellular hotspot is often IPv6-first under
244 # NAT64.
245 echo "n.add State:/Network/Global/IPv4"
246 echo "n.add State:/Network/Global/IPv6"
247 echo "n.watch"
248 while true; do sleep 86400; done
249 ) | script -q /dev/null scutil 2>/dev/null \
250 | while IFS= read -r line; do
251 case "$line" in
252 *changedKey*|*State:/Network/Global/IPv*|*n.state*|*SCEventUpdate*)
253 detect_lan_p2p_mode
254 run_recover "NET: $(echo "$line" | head -c 100)"
255 ;;
256 *) ;;
257 esac
258 done
259 echo "[$(date -u +%FT%TZ)] scutil pipe exited; reconnecting in 5s"
260 sleep 5
261 done
262 ) &
263 # Listener 3: TUNNEL_FAILED_CLOSED.marker Watcher
264 # Polls for the fail-closed marker dropped by the routing-fix container.
265 (
266 last_notified=0
267 while ;; do
268 if [ -f "$SCRIPT_DIR/../TUNNEL_FAILED_CLOSED.marker" ]; then
269 now=$(date +%s)
270 # Notify once every 5 minutes (300s) if it stays failed
271 if [ "$(now - last_notified)" -ge 300 ]; then
272 osascript -e 'display notification "VPN Tunnel Health Check Failed. Internet traffic is
blackholed to prevent IP leak." with title "NullExit Alert"' || true
273 last_notified=$now
274 fi
275 else
276 last_notified=0
277 fi
278 sleep 3
279 done
280 ) &
281
282 # Lifetime
283 # 'wait' blocks until both background pipelines exit (which they normally
284 # never do). launchctl 'bootout'/SIGTERM triggers the trap, which kills the
285 # whole process group so the sleep-forever in the scutil pipe also dies.
286 # We catch TERM/INT/HUP for explicit signals and EXIT for any normal-exit path
287 # (so listeners always get cleaned up even if 'exit 0' runs somewhere).
288 # 'rm -f $LOCK_FILE' first so a stale lock can never block the next launch.
289 # 'pkill -TERM -P $$' then targets direct listener-children, then 'kill 0'
290 # takes care of any backgrounded group-mates not reachable via the parent.
291 trap 'echo "[$(date -u +%FT%TZ)] watcher.sh terminating (signal/exit=$?)"; rm -f "$LOCK_FILE" 2>/dev/null
|| true; pkill -TERM -P $$ 2>/dev/null || true; kill 0 2>/dev/null || true; exit 0' TERM INT HUP
EXIT
292 wait

```

13 scripts/crypto.sh

```

1 #!/usr/bin/env bash
2 # crypto.sh - Cryptographic Integrity Check
3
4 set -euo pipefail
5 cd "$(dirname "$0")/.."

```

```

6
7 if [ "$#" -ne 1 ] || { [ "$1" != "--sign" ] && [ "$1" != "--verify" ]; }; then
8     echo "Usage: $0 --sign | --verify"
9     exit 1
10 fi
11
12 if [ ! -f .env ]; then
13     echo "Error: .env not found."
14     exit 1
15 fi
16
17 SEED=$(grep "^NULLEXIT_SEED=" .env | cut -d '=' -f2)
18 if [ -z "$SEED" ]; then
19     echo "Error: NULLEXIT_SEED not found in .env."
20     exit 1
21 fi
22
23 if [ "$1" == "--sign" ]; then
24     echo "Generating cryptographic signatures using seed..."
25     rm -f .signatures
26     for file in toggle.sh recover.sh scripts/common.sh scripts/routing-fix.sh scripts/diagnose-host-leak.sh
27         scripts/sweep.sh scripts/noise.sh scripts/watcher.sh scripts/pf.conf post-rules.txt docker-compose.
28         yml black_list.txt white_list.txt; do
29         if [ -f "$file" ]; then
30             hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
31             echo "$file:$hash" >> .signatures
32             echo " [Signed] $file"
33         else
34             echo " [Skipped] $file not found"
35         fi
36     done
37     echo "Signatures saved to .signatures"
38     exit 0
39 fi
40
41 if [ "$1" == "--verify" ]; then
42     if [ ! -f .signatures ]; then
43         echo "Error: .signatures file missing. Run $0 --sign first."
44         exit 1
45     fi
46     FAIL=0
47     while IFS=: read -r file expected_hash; do
48         if [ -z "$file" ]; then continue; fi
49         if [ ! -f "$file" ]; then
50             echo "[FAIL] $file is missing!"
51             FAIL=1
52             continue
53         fi
54         actual_hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
55         if [ "$actual_hash" != "$expected_hash" ]; then
56             echo "[FAIL] $file has been modified! Hash mismatch."
57             FAIL=1
58         fi
59     done < .signatures
60
61     if [ "$FAIL" -eq 1 ]; then
62         echo ""
63         echo "CRITICAL: Script integrity verification failed!"
64         echo "One or more core files have been modified without authorization."
65         echo "To authorize these changes, run: ./scripts/crypto.sh --sign"
66         echo ""
67         exit 1
68     fi
69     exit 0
70 fi

```

14 .aiignore

```

1 # Environment variables and secrets
2 .env
3 tor-bridges.txt

```

15 .claude/scheduled_tasks.lock

```
1 {"sessionId":"43bd1397-8ada-422e-9c6a-2f5c0ee5c708","pid":47318,"procStart":"Sat Jul 18 14:23:53 2026","acquiredAt":1784388323901}
```

16 .claude/settings.local.json

```
1 {
2   "permissions": {
3     "allow": [
4       "Bash(grep *)",
5       "Bash(python3 /tmp/gen_nullexit_graph.py)",
6       "Bash(python3 -)",
7       "Bash(`/Applications/Google Chrome.app/Contents/MacOS/Google Chrome` --headless=new --disable-gpu --window-size=1280,800 --virtual-time-budget=8000 --dump-dom `file:///tmp/kg_debug.html`)",
8       "Bash(echo `dangling-check-exit=$?`)"
9     ]
10  }
11 }
```

17 .env.example

```
1 # nullexit configuration template copy to .env and fill in.
2 # cp .env.example .env
3 # .env is gitignored (it holds secrets); this template documents every option
4 # with its shipping default. For coherent bundles, see scripts/profiles.sh.
5
6 # Secrets (fill these in)
7 # WARP WireGuard keys from your Cloudflare WARP registration.
8 WIREGUARD_PRIVATE_KEY=<your-warp-wireguard-private-key>
9 WIREGUARD_PUBLIC_KEY=<your-warp-wireguard-public-key>
10 WIREGUARD_ADDRESSES=172.16.0.2/32
11 # Tailscale auth key Admin Console Settings Keys Generate auth key.
12 TS_AUTHKEY=tskey-auth-xxxxxxxxxxxxxxxxxxxxxxxx
13 # HMAC seed for script-integrity signatures generate with: openssl rand -hex 32
14 NULLEXIT_SEED=<run: openssl rand -hex 32>
15 # AdGuard Home UI + Tor ControlPort credentials.
16 ADGUARD_USER=admin
17 ADGUARD_PASSWORD=<choose-a-password>
18
19 # Gateway core
20 # Rule-compiler memory profile: light (~52k) / medium (~167k) / heavy (~253k).
21 GATEWAY_RULE_PROFILE=medium
22 STOP_COLIMA_ON_EXIT=true
23 # PF fail-closed kill-switch. Leave true this is the core leak guarantee.
24 KILL_SWITCH=true
25 # true only on trusted home networks (no AP isolation); false on campus/enterprise
26 # WPA2-Enterprise (forces DERP, avoids gvproxy SNAT poisoning).
27 TAILSCALE_ALLOW_LAN_P2P=false
28 GATEWAY_MSS=1120
29 GATEWAY_BYPASS_PING=false
30 GATEWAY_USE_EXIT_NODE=true
31 # Consecutive warp=off polls before auto-shutdown (6 = 30s; 3 = faster/stricter).
32 WARP_FAIL_THRESHOLD=6
33 BLOCKED_COUNTRIES="kp il"
34 HOST_MTU=1200
35 # Debug xtrace of toggle.sh to output.log (keep false for normal use).
36 DEBUG_TRACE=false
37
38 # Tor hardening
39 # true = Tor must use obfs4 bridges only (never public guards).
40 TOR_USE_BRIDGES=false
41 TOR_BRIDGE_LINES_FILE=./tor-bridges.txt
42 TOR_EXCLUDE_EXIT_NODES={us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il}
43
44 # Cover traffic / dummy padding (scripts/noise.sh) opt-in, default off
45 NOISE_ENABLED=false
46 NOISE_PAD_TARGET=
47 NOISE_PAD_PORT=9999
48 NOISE_PAD_RATE_KBPS=64
49 NOISE_PAD_PACKET_BYTES=1024
```

```

50 NOISE_PAD_JITTER_MS=40
51 # constant = always emit the target rate; topup = fill real-traffic gaps to keep
52 # the observed rate flat (stronger anti-correlation, no waste when already busy).
53 NOISE_PAD_MODE=constant
54 NOISE_PAD_IFACE=en0
55 # Optional: target as a % of LINK CAPACITY (overrides NOISE_PAD_RATE_KBPS).
56 NOISE_PAD_PCT=
57 NOISE_LINK_MBPS=100
58
59 # FaceTime split-route (scripts/common.sh) opt-in, default off
60 # true routes the Apple /16s below DIRECT on your REAL IP (not WARP) so
61 # FaceTime can hole-punch. A deliberate, narrow hole scoped to Apple infra.
62 FACETIME_SPLIT_ROUTE=false
63 FACETIME_DIRECT_SUBNETS="17.249.0.0/16"
64
65 # LAN Pi-hole mode (scripts/pihole-mode.sh) opt-in, default off
66 # Serve AdGuard DNS filtering to non-Tailscale LAN devices (trusted networks only).
67 PIHOLE_LAN_MODE=false
68 PIHOLE_LAN_LISTEN=
69
70 # Sweep throughput check (scripts/sweep.sh, check 7/7) all optional
71 # Uncomment to override the defaults baked into sweep.sh (a fast CDN, 30s cap).
72 # SWEEP_THROUGHPUT_URL=https://ash-speed.hetzner.com/100MB.bin
73 # SWEEP_THROUGHPUT_TIMEOUT=30
74 # SWEEP_SKIP_THROUGHPUT=false

```

18 .gitignore

```

1 # Environment variables and secrets
2 .env
3
4 # wgcf generated files containing keys and account info
5 wgcf-account.toml
6 wgcf-profile.conf
7
8 # Tailscale persistent state containing node identity
9 tailscale-state/
10
11 # Binaries
12 wgcf
13
14 # macOS App bundles
15 *.app/
16 *.app
17
18 # AdGuard Home persistent state
19 adguard/
20 .gateway_ip
21 logs.txt
22
23 #
24 stability_incident_report.md
25 output.log
26 .DS_Store
27 *.desktop
28 blocked.log
29 wtf.md
30 host-leak-diagnostic-*.txt
31 nullexit_review.md
32 nullexit_unified.tex
33 nullexit_unified.txt
34
35 nullexit_unified.synctex.gz
36 tor-bridges.txt
37
38 # Runtime state files
39 .host_ips
40 .lan_p2p_detected
41 .build_hash
42 .rules_hash
43 .signatures
44 TUNNEL_FAILED_CLOSED.marker
45
46 # Python cache
47 __pycache__/_
48 *.py[cod]

```

```

49
50 # LaTeX build artifacts
51 *.aux
52 *.log
53 *.out
54 *.toc
55 *.xdv
56 *.fdb_latexmk
57 *.fls
58 *.synctex*
59 refactor.md
60 sweep.md
61
62 # Obsidian Knowledge graph
63 nullexit-knowledge-graph/
64
65 # DNS anomaly detector per-network model (regenerated by 'learn')
66 .dns_baseline.json
67
68 # .env backups written by profiles.sh (contain secrets never commit)
69 .env.bak*
70
71 # device-scramble saved real identity (never commit)
72 .device-identity.orig

```

19 Recover-Gateway.applescript

```

1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3     do shell script "true" with administrator privileges
4 on error
5     return
6 end try
7 tell application "Terminal"
8     activate
9     do script "cd " & quoted form of scriptDir & "; bash recover.sh"
10 end tell

```

20 Sweep-Gateway.applescript

```

1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3     do shell script "cd " & quoted form of scriptDir & " && ./scripts/crypto.sh --verify"
4 on error errMsg
5     display dialog "nullexit integrity check FAILED sweep will not run." & return & return & "A signed
        file (scripts/sweep.sh, scripts/common.sh, ) no longer matches .signatures. If you edited it on
        purpose, re-sign with:" & return & ".scripts/crypto.sh --sign" & return & return & errMsg buttons {
            "Abort"} default button "Abort" with icon stop
6     return
7 end try
8 tell application "Terminal"
9     activate
10    do script "cd " & quoted form of scriptDir & " && bash scripts/sweep.sh"
11 end tell

```

21 Toggle-Gateway.applescript

```

1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3     do shell script "true" with administrator privileges
4 on error
5     return
6 end try
7 tell application "Terminal"
8     activate
9     do script "cd " & quoted form of scriptDir & "; ./toggle.sh"
10 end tell

```

22 benchmarks/benchmark.sh

```
1 #!/usr/bin/env bash
2
3 # benchmark.sh
4 # Automates the full benchmarking suite (both Python download test and Lighthouse HTML test)
5 # Runs the suite with Gateway ON, toggles OFF, runs again, then toggles back ON.
6
7 set -e
8
9 # Change to project root so we can access toggle.sh reliably
10 cd "$(dirname "$0")/.."
11
12 # Ensure jq is installed
13 if ! command -v jq &> /dev/null; then
14     echo "Error: 'jq' is not installed. Please run 'brew install jq'"
15     exit 1
16 fi
17
18 run_lighthouse() {
19     local url=$1
20     local file_prefix=$2
21     echo " [Lighthouse] Testing $url ..."
22
23     # Run lighthouse headlessly, silencing standard output
24     npx -y lighthouse "$url" --chrome-flags="--headless" --only-categories=performance --output=json --
25     output-path="${file_prefix}.json" >/dev/null 2>&1
26
27     # Parse the results
28     local score=$(jq -r '.categories.performance.score' "${file_prefix}.json")
29     local tti=$(jq -r '.audits["interactive"].displayValue' "${file_prefix}.json")
30     local si=$(jq -r '.audits["speed-index"].displayValue' "${file_prefix}.json")
31
32     # Convert score to percentage
33     local score_pct=$(awk "BEGIN {print $score*100}")
34
35     echo "    -> Performance Score: ${score_pct}%"
36     echo "    -> Time to Interactive: $tti"
37     echo "    -> Speed Index: $si"
38
39     # Cleanup
40     rm -f "${file_prefix}.json"
41 }
42
43 echo "=====
44 echo " PHASE 1: Testing with Gateway ON"
45 echo "=====
46 # 1. Run raw Python download test
47 python3 benchmarks/test_load.py
48
49 # 2. Run Lighthouse HTML render test
50 echo ""
51 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
52 run_lighthouse "https://www.cnet.com/" "on_cnet"
53 run_lighthouse "https://www.independent.co.uk/" "on_ind"
54
55 echo ""
56 echo "=====
57 echo " PHASE 2: Toggling Gateway OFF"
58 echo "=====
59 ./toggle.sh
60 echo "Waiting 10 seconds for macOS Wi-Fi to stabilize after DNS reset..."
61 sleep 10
62
63
64 echo ""
65 echo "=====
66 echo " PHASE 3: Testing with Gateway OFF"
67 echo "=====
68 # 1. Run raw Python download test
69 python3 benchmarks/test_load.py
70
71 # 2. Run Lighthouse HTML render test
72 echo ""
73 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
74 run_lighthouse "https://www.cnet.com/" "off_cnet"
75 run_lighthouse "https://www.independent.co.uk/" "off_ind"
76
```



```

77
78 echo ""
79 echo "=====
80 echo "    PHASE 4: Restoring Gateway ON"
81 echo "=====
82 ./toggle.sh
83
84 echo ""
85 echo "=====
86 echo " BENCHMARKS COMPLETE!"
87 echo "You can scroll up to compare Phase 1 (AdGuard ON) vs Phase 3 (AdGuard OFF)."
88 echo "=====

```

23 benchmarks/test_load.py

```

1 import urllib.request
2 import urllib.error
3 import time
4 from html.parser import HTMLParser
5 import concurrent.futures
6
7 class ResourceParser(HTMLParser):
8     def __init__(self):
9         super().__init__()
10        self.urls = set()
11
12    def handle_starttag(self, tag, attrs):
13        attrs = dict(attrs)
14        url = None
15        if tag in ['img', 'script']:
16            url = attrs.get('src')
17        elif tag == 'link' and attrs.get('rel') == 'stylesheet':
18            url = attrs.get('href')
19
20        if url and url.startswith('http'):
21            self.urls.add(url)
22
23 def fetch_url(url):
24     try:
25         req = urllib.request.Request(url, headers={'User-Agent': 'Mozilla/5.0'})
26         urllib.request.urlopen(req, timeout=3)
27         return True
28     except Exception:
29         return False
30
31 def measure_url(target_url):
32     print(f"\nTesting {target_url} Load...")
33     start_total = time.time()
34
35     # Fetch main HTML
36     try:
37         req = urllib.request.Request(target_url, headers={'User-Agent': 'Mozilla/5.0 (Macintosh; Intel
38         Mac OS X 10_15_7) AppleWebKit/537.36'})
39         response = urllib.request.urlopen(req, timeout=5)
40         html = response.read().decode('utf-8', errors='ignore')
41     except Exception as e:
42         print(f"Failed to fetch {target_url}: {e}")
43         return
44
45     parser = ResourceParser()
46     parser.feed(html)
47     urls = list(parser.urls)
48     print(f"Found {len(urls)} subresources to fetch (scripts, images, css)...")
49
50     # Concurrently fetch subresources
51     success_count = 0
52     fail_count = 0
53     max_threads = min(64, max(1, len(urls)))
54     with concurrent.futures.ThreadPoolExecutor(max_workers=max_threads) as executor:
55         results = list(executor.map(fetch_url, urls))
56
57     success_count = sum(1 for r in results if r)
58     fail_count = len(results) - success_count
59
60     total_time = time.time() - start_total
61     print(f"Time taken: {total_time:.2f} seconds")

```

```

61     print(f"Successfully fetched: {success_count}, Failed/Blocked: {fail_count}")
62
63     urls_to_test = [
64         "https://www.dailymail.co.uk/home/index.html",
65         "https://www.cnet.com/",
66         "https://www.independent.co.uk/"
67     ]
68
69     for u in urls_to_test:
70         measure_url(u)

```

24 black_list.txt

```

1  ad.doubleclick.net
2  adclick.g.doubleclick.net
3  adeventtracker.spotify.com
4  adfstat.yandex.ru
5  ads-api.tiktok.com
6  ads-api.twitter.com
7  ads-fa.spotify.com
8  ads-sg.tiktok.com
9  ads.tiktok.com
10 adserver.unityads.unity3d.com
11 adsfs.oppomobile.com
12 an.facebook.com$important
13 connect.facebook.net$important
14 pixel.facebook.com$important
15 analytics.query.yahoo.com
16 analytics.s3.amazonaws.com
17 analytics.spotify.com
18 analyticsengine.s3.amazonaws.com
19 appmetrica.yandex.ru
20 audio2.spotify.com
21 bdapi-ads.realmemobile.com
22 bdapi-in-ads.realmemobile.com
23 bounceexchange.com
24 bs.serving-sys.com
25 business-api.tiktok.com
26 ck.ads.oppomobile.com
27 crashdump.spotify.com
28 data.ads.oppomobile.com
29 data.mistat.india.xiaomi.com
30 data.mistat.rus.xiaomi.com
31 data.mistat.xiaomi.com
32 desktop.spotify.com
33 doubleclick.net
34 ds.metric.spotify.com
35 gemini.yahoo.com
36 googleads.g.doubleclick.net
37 googleadservices.com
38 googlesyndication.com
39 grs.hicloud.com
40 gtm.spotify.com
41 iot-eu-logser.realme.com
42 iot-logser.realme.com
43 log.byteoversea.com
44 log.fc.yahoo.com
45 log.spotify.com
46 m.doubleclick.net
47 mediavisor.doubleclick.net
48 metrics.spotify.com
49 metrika.yandex.ru
50 omaze.com
51 pagead2.googlesyndication.com
52 securepubads.g.doubleclick.net
53 static.doubleclick.net
54 stats.g.doubleclick.net
55 tracking.rus.miui.com
56 udcms.yahoo.com
57 v.jwpcdn.com
58 weblb-wg.gslb.spotify.com
59 webview.unityads.unity3d.com
60 diagassets.apple.com
61 iphonesubmissions.apple.com
62 radarsubmissions.apple.com
63 metrics.apple.com

```

```

64 xp.apple.com
65 heads-fa.spotify.com
66 heads4.spotify.com
67 pixel.spotify.com
68 segs.spotify.com
69 audio-fa.spotify.com
70 events.data.microsoft.com
71 datarouter-prod.ak.epicgames.com
72 nel.cloudflare.com
73 eu-mobile.events.data.microsoft.com
74 stocks-analytics-events.apple.com

```

25 cloud-init.yaml

```

1 #cloud-config
2 # Generic Cloud-Init Bootstrap Script for Remote Exit Nodes (Ubuntu/Debian)
3 # Paste this into the "User Data" or "Initialization Script" section when creating a VPS.
4
5 package_update: true
6 package_upgrade: true
7
8 packages:
9   - apt-transport-https
10  - ca-certificates
11  - curl
12  - gnupg
13  - lsb-release
14  - git
15  - python3
16  - python3-pip
17
18 runcmd:
19   # 1. Install Docker natively
20   - curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/
     docker-archive-keyring.gpg
21   - echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg
     ] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.
     list.d/docker.list > /dev/null
22   - sudo apt-get update
23   - sudo apt-get install -y docker-ce docker-ce-cli containerd.io docker-compose-plugin
24
25   # 2. Add default user (ubuntu) to docker group
26   - usermod -aG docker ubuntu || true
27   - usermod -aG docker debian || true
28
29   # 3. Create the installation directory
30   - mkdir -p /opt/nullexit
31   - chown -R ubuntu:ubuntu /opt/nullexit || chown -R debian:debian /opt/nullexit
32
33   # 4. Enable IPv4 and IPv6 forwarding for VPN routing
34   - echo "net.ipv4.ip_forward=1" >> /etc/sysctl.d/99-tailscale.conf
35   - echo "net.ipv6.conf.all.forwarding=1" >> /etc/sysctl.d/99-tailscale.conf
36   - sysctl -p /etc/sysctl.d/99-tailscale.conf
37
38 final_message: "The remote exit node has been fully bootstrapped. SSH in, clone the repo to /opt/nullexit
, add your .env, and run ./setup-linux.sh!"

```

26 docker/tor/Dockerfile

```

1 FROM debian:bookworm-slim
2
3 RUN apt-get update && \
4     apt-get install -y tor obfs4proxy gosu bash && \
5     rm -rf /var/lib/apt/lists/*
6
7 COPY entrypoint.sh /entrypoint.sh
8 RUN chmod +x /entrypoint.sh
9
10 ENTRYPOINT ["/entrypoint.sh"]

```

27 docker/tor/entrypoint.sh

```
1 #!/bin/bash
2 set -e
3
4 TORRC="/etc/tor/torrc"
5 mkdir -p /etc/tor /var/lib/tor
6
7 # Base configuration: SOCKS 9050, Control 9051, TransPort 9040, DNSPort 5353 all bound to
8 # 0.0.0.0 (required for Docker/Colima host port-mapping). The ControlPort is secured via a
9 # dynamic HashedControlPassword (added below when TOR_PASSWORD is set); see devref.md 15.10.2.
10 cat <<EOF > $TORRC
11 SocksPort 0.0.0.0:9050
12 ControlPort 0.0.0.0:9051
13 CookieAuthentication 0
14 Log notice stdout
15 DataDirectory /var/lib/tor
16
17 # Transparent proxying & DNS resolution for .onion mapping
18 TransPort 0.0.0.0:9040
19 DNSPort 0.0.0.0:5353
20 AutomapHostsOnResolve 1
21 VirtualAddrNetworkIPv4 198.18.0.0/15
22 EOF
23
24 if [ -n "$TOR_PASSWORD" ]; then
25     HASH=$(tor --hash-password "$TOR_PASSWORD" | tail -n 1)
26     echo "HashedControlPassword $HASH" >> $TORRC
27 fi
28
29 if [ -n "$TOR_EXCLUDE_EXIT_NODES" ]; then
30     echo "ExcludeExitNodes $TOR_EXCLUDE_EXIT_NODES" >> $TORRC
31     echo "StrictNodes 1" >> $TORRC
32 fi
33
34 if [ "${TOR_USE_BRIDGES:-false}" = "true" ]; then
35     BRIDGE_FILE="${TOR_BRIDGE_LINES_FILE:-/tor-bridges.txt}"
36
37     if ! grep -vE "^[[:space:]]*#" "$BRIDGE_FILE" | grep -q '^[[:space:]]'; then
38         MSG="[(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL] [Tor Proxy] TOR_USE_BRIDGES is enabled, but no
39         bridge lines found at $BRIDGE_FILE. Proxy startup ABORTED. Get bridges from https://bridges.
40         torproject.org."
41         echo "$MSG"
42         if [ -w "/output.log" ]; then
43             echo "$MSG" >> /output.log
44         fi
45         # Sleep indefinitely to prevent docker-compose 'unless-stopped' from triggering a crash loop
46         exec sleep infinity
47     fi
48
49     echo "UseBridges 1" >> $TORRC
50     echo "ClientTransportPlugin obfs4 exec /usr/bin/obfs4proxy" >> $TORRC
51
52     while IFS= read -r line; do
53         # Ignore comments and empty lines
54         [[ -z "$line" || "$line" == \#* ]] && continue
55         echo "Bridge $line" >> $TORRC
56     done < "$BRIDGE_FILE"
57 fi
58
59 # Debian tor package uses 'debian-tor' user
60 chown -R debian-tor:debian-tor /var/lib/tor /etc/tor
61
62 exec gosu debian-tor /usr/bin/tor -f $TORRC
```

28 launchd/com.nullexit.noise.plist

```
1 <?xml version="1.0" encoding="UTF-8"?>
2 <!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
3 <plist version="1.0">
4 <dict>
5     <key>Label</key>
6     <string>com.nullexit.noise</string>
7
8     <!-- Same install-agnostic pattern as com.nullexit.wake-recovery: WorkingDirectory
```

```

9      is __NULLEXIT_HOME__ (substitute your absolute install root at install time
10      via sed), and the program arg is scripts/noise.sh RELATIVE to it, so
11      BASH_SOURCE still resolves SCRIPT_DIR correctly. -->
12      <key>WorkingDirectory</key>
13      <string>__NULLEXIT_HOME__</string>
14
15      <!-- ‘__boot’: governed SOLELY by NOISE_ENABLED in .env. It exits 0 cleanly when
16      disabled (so KeepAlive never tight-loops) and runs the padding loop in the
17      foreground when enabled. Loading this plist adds NO new option it just means
18      "auto-start cover traffic whenever NOISE_ENABLED=true". -->
19      <key>ProgramArguments</key>
20      <array>
21          <string>/bin/bash</string>
22          <string>scripts/noise.sh</string>
23          <string>__boot</string>
24      </array>
25
26      <!-- Start immediately on launchctl load and re-run at every login/boot. -->
27      <key>RunAtLoad</key>
28      <true/>
29
30      <!-- Same restart policy rationale as wake-recovery:
31      SuccessfulExit=false a clean exit 0 (the disabled-idle path) is NOT
32                          relaunched, so a disabled job doesn't spin.
33      Crashed=false       a SIGTERM from ‘launchctl bootout’ is not treated as
34                          a relaunched trigger, so stopping is clean and loop-free.
35      Net effect: load once padding runs at boot when enabled; bootout stops it. -->
36      <key>KeepAlive</key>
37      <dict>
38          <key>SuccessfulExit</key>
39          <false/>
40          <key>Crashed</key>
41          <false/>
42      </dict>
43
44      <!-- Cover traffic must not be suspended by App Nap when the user is idle. -->
45      <key>ProcessType</key>
46      <string>Background</string>
47
48      <key>StandardOutPath</key>
49      <string>/tmp/nullexit-noise.out.log</string>
50      <key>StandardErrorPath</key>
51      <string>/tmp/nullexit-noise.err.log</string>
52  </dict>
53 </plist>

```

29 launchd/com.nullexit.wake-recovery.plist

```

1  <?xml version="1.0" encoding="UTF-8"?>
2  <!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
3  <plist version="1.0">
4  <dict>
5      <key>Label</key>
6      <string>com.nullexit.wake-recovery</string>
7
8      <!-- WorkingDirectory + relative program arg keeps this plist install-agnostic.
9      The source-of-truth for the install path is __NULLEXIT_HOME__, which
10     setup.sh or the user must substitute with their absolute nullexit
11     install root at install time. Program arg is ‘scripts/watcher.sh’
12     RELATIVE to WorkingDirectory, so launchd’s evaluated cwd IS the
13     install root and BASH_SOURCE inside scripts/watcher.sh still
14     resolves to the real <install>/scripts/watcher.sh, preserving the
15     SCRIPT_DIR pattern. After install-time ‘sed -i ’ ’ ’s|__NULLEXIT_HOME__|<abs_path>|’ ’,
16     the plist is portable to any clone location. -->
17     <key>WorkingDirectory</key>
18     <string>__NULLEXIT_HOME__</string>
19
20     <key>ProgramArguments</key>
21     <array>
22         <string>/bin/bash</string>
23         <string>scripts/watcher.sh</string>
24     </array>
25
26     <!-- Start the daemon immediately on launchctl load (instead of waiting
27     for the next user login). Combined with keepAlive below this means
28     once installed: load once, runs forever across logouts/reboots. -->

```

```

29     <key>RunAtLoad</key>
30     <true/>
31
32     <!-- Restart policy:
33         SuccessfulExit=false    SIGTERM / clean exit does NOT relaunch (so
34         'launchctl bootout' cleanly stops, not loops).
35         Crashed=false          HARD crashes ALSO do NOT relaunch. We've seen
36                                that repeated sleep/wake cycles on macOS
37                                Sonoma+ accumulate orphan watcher.sh PIDs
38                                because SIGCONT/resume doesn't reliably kill
39                                the prior instance. The watcher.sh script
40                                now enforces single-instance via a manual
41                                PID-file and process-identity check at the
42                                top, so a re-launch on crash is
43                                actively dangerous (it would skip the lock
44                                and exit; better to surface the crash in
45                                logs than to silently 0-process). -->
46
47     <key>KeepAlive</key>
48     <dict>
49         <key>SuccessfulExit</key>
50         <false/>
51         <key>Crashed</key>
52         <false/>
53     </dict>
54
55     <!-- Background hint to App Nap so it doesn't get suspended when
56         the user is inactive. We're the entire reason this LaunchAgent
57         needs to NOT nap. -->
58     <key>ProcessType</key>
59     <string>Background</string>
60
61     <!-- launchd-level stdout/stderr (the bash launcher shell). The script
62         itself writes to /tmp/nullexit-watcher.log via exec >>. These
63         separate channels help us tell "launchd crashed" from "the bash
64         wrapper piped to recover.sh errored". -->
65     <key>StandardOutPath</key>
66     <string>/tmp/nullexit-watcher.out.log</string>
67     <key>StandardErrorPath</key>
68     <string>/tmp/nullexit-watcher.err.log</string>
69 </dict>
</plist>

```

30 post-rules.txt

```

1 # NOTE: the REJECT rules below (DNS-over-TLS, QUIC) must be added BEFORE the blanket
2 # FORWARD ACCEPT rules. iptables stops at the first match, so if the ACCEPT rules ran
3 # first, all tailscale0 traffic (including port 853/443) would already be accepted and
4 # these REJECT rules would never be reached they'd sit dead with a permanent 0-packet counter.
5
6 # Block DNS-over-TLS (Port 853) to force Android Private DNS to fallback to Port 53
7 iptables -A FORWARD -i tailscale0 -p tcp --dport 853 -j REJECT --reject-with tcp-reset
8 iptables -A FORWARD -i tailscale0 -p udp --dport 853 -j REJECT
9
10 # Block QUIC (UDP 443) to force Facebook/Google to fallback to TCP.
11 # UDP bypasses TCP MSS clamping, causing fragmentation. Over weak Wi-Fi (far from gateway),
12 # dropping a single fragment destroys the entire packet, stalling image downloads.
13 iptables -A FORWARD -i tailscale0 -p udp --dport 443 -j REJECT
14
15 iptables -A FORWARD -i tailscale0 -o tun0 -j ACCEPT
16 iptables -A FORWARD -i tun0 -o tailscale0 -j ACCEPT
17 iptables -A INPUT -i tailscale0 -j ACCEPT
18 iptables -A INPUT -i eth0 -p udp --dport 41642 -j ACCEPT
19 iptables -t nat -A POSTROUTING -o tun0 -j MASQUERADE
20 # Clamp MSS to 1120. With Tailscale overhead on top of WARP (each WireGuard encapsulation adds ~60 bytes)
21 # the safe MSS for double-tunneled traffic is 1120 to prevent strict-MSS endpoints from stalling.
22 # NOTE: this 1120 integer is rewritten at boot by toggle.sh from GATEWAY_MSS in .env
23 # (Gluetun's parser can't interpolate variables) change GATEWAY_MSS in .env, not this line.
24 iptables -t mangle -A POSTROUTING -p tcp --tcp-flags SYN,RST SYN -j TCPMSS --set-mss 1120
25
26 iptables -t nat -A PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT --to-port 5335
27 iptables -t nat -A PREROUTING -i tailscale0 -p tcp --dport 53 -j REDIRECT --to-port 5335
28
29 # Also redirect DNS from Docker bridge interface so the host can reach AdGuard via localhost:53
30 iptables -t nat -A PREROUTING -i eth0 -p udp --dport 53 -j REDIRECT --to-port 5335
31 iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 53 -j REDIRECT --to-port 5335

```

```

32
33 # Block IPv6 exit-node traffic to prevent leakage (Gluetun/WARP is IPv4-only)
34 iptables -A FORWARD -i tailscale0 -j DROP

```

31 scripts/Dockerfile.routing-fix

```

1 FROM golang:1.22-alpine AS builder
2 WORKDIR /app
3 COPY logger/ ./logger/
4 RUN cd logger && go build -o logger main.go
5
6 FROM alpine:3.20
7 RUN apk add --no-cache iptables iproute2 curl ipset
8 COPY --from=builder /app/logger/logger /usr/local/bin/nullexit-logger
9 COPY routing-fix.sh /routing-fix.sh
10 CMD ["sh", "/routing-fix.sh"]

```

32 scripts/Dockerfile.rule-compiler

```

1 FROM golang:1.22-alpine AS builder
2 WORKDIR /app
3 COPY rule-compiler/ ./rule-compiler/
4 RUN cd rule-compiler && go build -o sync-rules main.go
5
6 FROM alpine:3.20
7 COPY --from=builder /app/rule-compiler/sync-rules /usr/local/bin/sync-rules
8 WORKDIR /app
9 CMD ["sync-rules"]

```

33 scripts/Toggle-Gateway.bat

```

1 <# :
2 @echo off
3 powershell -NoProfile -ExecutionPolicy Bypass -Command "Invoke-Expression ([System.IO.File]::ReadAllText
4 ('%f0') -replace '^(?s).*?#\s*>\s*\r?\n', '')"
5 exit /b
6 #>
7 # Toggle-Gateway.bat (Hybrid Batch-PowerShell Toggler) nullexit Gateway Toggler for Windows
8 # Mirrors the sophisticated error handling, checks, and automation of toggle.sh
9
10 # 1. Administrator Check & Elevation
11 $isAdmin = ([Security.Principal.WindowsPrincipal][Security.Principal.WindowsIdentity]::GetCurrent()).
12 IsInRole([Security.Principal.WindowsBuiltInRole]::Administrator)
13 if (-not $isAdmin) {
14     Write-Host "Requesting Administrator privileges..." -ForegroundColor Yellow
15     Start-Process PowerShell -ArgumentList "-NoProfile -ExecutionPolicy Bypass -Command 'Invoke-
16 Expression ([System.IO.File]::ReadAllText('PSCommandPath') -replace '^(?s).*?#\s*>\s*\r?\n', '')'"
17     -Verb RunAs
18     Exit
19 }
20
21 # Set console output encoding to UTF8 for clean symbols
22 [Console]::OutputEncoding = [System.Text.Encoding]::UTF8
23 Clear-Host
24
25 Write-Host "" -ForegroundColor Green
26 Write-Host " nullexit Windows Toggler " -ForegroundColor Green
27 Write-Host "" -ForegroundColor Green
28 Write-Host ""
29
30 # 2. Helpers
31 function Reset-HostDNS {
32     Write-Host "Restoring default DNS (DHCP/DHCP-assigned DNS)..." -ForegroundColor Cyan
33     $ActiveAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
34     foreach ($adapter in $ActiveAdapters) {
35         Write-Host " -> Resetting DNS on adapter: $($adapter.Name)" -ForegroundColor Gray
36         Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ResetServerAddresses -
37         ErrorAction SilentlyContinue
38     }
39 }

```



```

33     }
34 }
35
36 function Hijack-HostDNS ($ip) {
37     Write-Host "Hijacking DNS to route through AdGuard ($ip)..." -ForegroundColor Yellow
38     $activeAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
39     foreach ($adapter in $activeAdapters) {
40         Write-Host "    -> Pointing DNS to $ip on adapter: $($adapter.Name)" -ForegroundColor Gray
41         Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ServerAddresses $ip -
            ErrorAction SilentlyContinue
42     }
43 }
44
45 function Get-HostTailscalePath {
46     $paths = @(
47         "$env:ProgramFiles\Tailscale\tailscale.exe",
48         "$env:ProgramFiles (x86)\Tailscale\tailscale.exe",
49         "tailscale.exe"
50     )
51     foreach ($path in $paths) {
52         if (Test-Path $path) { return $path }
53     }
54     return $null
55 }
56
57 # 3. Detect State
58 # Prevent deadlocks by resetting DNS to default temporarily during checks
59 Reset-HostDNS
60
61 Write-Host "Checking Docker daemon status..." -ForegroundColor Cyan
62 & docker info >$null 2>&1
63 if ($LASTEXITCODE -ne 0) {
64     Write-Host "Docker is not running. Starting Docker Desktop..." -ForegroundColor Yellow
65     $dockerPath = "C:\Program Files\Docker\Docker\Docker Desktop.exe"
66     if (Test-Path $dockerPath) {
67         Start-Process $dockerPath
68         Write-Host "Waiting for Docker daemon to become responsive..." -ForegroundColor Gray
69         $timeout = 60
70         while ($timeout -gt 0) {
71             Start-Sleep -Seconds 2
72             & docker info >$null 2>&1
73             if ($LASTEXITCODE -eq 0) {
74                 Write-Host "Docker is ready!" -ForegroundColor Green
75                 break
76             }
77             $timeout -= 2
78         }
79         if ($timeout -le 0) {
80             Write-Error "Docker failed to start in time. Aborting."
81             Exit 1
82         }
83     } else {
84         Write-Error "Docker Desktop executable not found at standard path. Please start Docker manually."
85         Exit 1
86     }
87 }
88
89 # Detect if gateway is already running
90 $runningWarp = docker compose ps --status running --format json | ConvertFrom-Json -ErrorAction
    SilentlyContinue | Where-Object { $_.Service -eq "warp" }
91 $isGatewayActive = ($null -ne $runningWarp)
92
93 # 4. Execution
94 if ($isGatewayActive) {
95     # STOP PATH
96     Write-Host "Gateway is currently RUNNING. Disabling now..." -ForegroundColor Yellow
97     Write-Host ""
98
99     # Disconnect host Tailscale from exit node if installed
100     $tsCli = Get-HostTailscalePath
101     if ($null -ne $tsCli) {
102         Write-Host "Disconnecting host Tailscale from exit node..." -ForegroundColor Cyan
103         & $tsCli up --exit-node= --accept-dns=true >$null 2>&1
104     }
105
106     Write-Host "Stopping Docker containers..." -ForegroundColor Cyan
107     docker compose down -t 5
108
109     Reset-HostDNS
110     Write-Host "Gateway has been successfully STOPPED." -ForegroundColor Green
111 } else {

```

```

112 # START PATH
113 Write-Host "Gateway is currently STOPPED. Enabling now..." -ForegroundColor Yellow
114 Write-Host ""
115
116 # Clean up corrupted AdGuardHome configurations
117 $confFile = "adguard\conf\AdGuardHome.yaml"
118 if (Test-Path $confFile) {
119     $fileSize = (Get-Item $confFile).Length
120     if ($fileSize -eq 0) {
121         Remove-Item $confFile -Force
122         Write-Host "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop." -
123             ForegroundColor Yellow
124     }
125 }
126
127 Write-Host "Starting Docker containers..." -ForegroundColor Cyan
128 docker compose up -d
129
130 Write-Host "Waiting for container's Tailscale connection to offer Exit Node..." -ForegroundColor Cyan
131 Write-Host " (This can take up to 60 seconds..." -ForegroundColor Gray
132
133 $elapsed = 0
134 $maxWait = 60
135 $resolvedIP = $null
136
137 while ($elapsed -lt $maxWait) {
138     $statusJson = docker compose exec -T tailscale tailscale status --json 2>$null
139     if ($null -ne $statusJson) {
140         $status = $statusJson | ConvertFrom-Json -ErrorAction SilentlyContinue
141         if ($null -ne $status -and $null -ne $status.Self) {
142             # Check if it has a valid Tailscale IP and is running
143             if ($status.Self.Online -and $status.Self.TailscaleIPs.Count -gt 0) {
144                 $resolvedIP = $status.Self.TailscaleIPs[0]
145                 break
146             }
147         }
148     }
149     Start-Sleep -Seconds 2
150     $elapsed += 2
151 }
152
153 if ($null -ne $resolvedIP) {
154     Write-Host "Tailscale IP resolved: $resolvedIP" -ForegroundColor Green
155     Write-Host ""
156
157     # Configure host Tailscale client to route through this exit node
158     $tsCli = Get-HostTailscalePath
159     if ($null -ne $tsCli) {
160         Write-Host "Configuring host Tailscale to use Exit Node ($resolvedIP)..." -ForegroundColor
161         Cyan
162         & $tsCli up --exit-node=$resolvedIP --accept-dns=false >$null 2>&1
163     }
164
165     # Hijack DNS
166     Hijack-HostDNS $resolvedIP
167     Write-Host ""
168     Write-Host "Gateway has been successfully STARTED." -ForegroundColor Green
169 } else {
170     Write-Host "Warning: Could not resolve gateway Tailscale IP. DNS hijacking skipped." -
171         ForegroundColor Red
172
173     # Fallback to .gateway_ip if it exists
174     if (Test-Path ".gateway_ip") {
175         $fallbackIP = Get-Content ".gateway_ip" -TotalCount 1
176         if (-not [string]::IsNullOrEmpty($fallbackIP)) {
177             Write-Host "[Fallback] Using static IP from .gateway_ip: $fallbackIP" -ForegroundColor
178             Yellow
179             Hijack-HostDNS $fallbackIP
180             Write-Host "Gateway STARTED (with static fallback IP)." -ForegroundColor Green
181         }
182     } else {
183         Write-Host "Gateway started in offline/unauthenticated state. Run setup.sh if this is a fresh
184         setup." -ForegroundColor Yellow
185     }
186 }
187
188 Write-Host ""
189 Write-Host "Press any key to exit..."
190 [void]$Host.UI.RawUI.ReadKey("NoEcho,IncludeKeyDown")

```

34 scripts/check_circuits.py

```
1 #!/usr/bin/env python3
2 import socket
3 import os
4 import re
5
6 def query_tor(port, command, password=""):
7     try:
8         s = socket.socket()
9         s.settimeout(2)
10        s.connect(('127.0.0.1', port))
11        s.sendall(f"AUTHENTICATE \"{password}\"\\r\\n{command}\\r\\nQUIT\\r\\n".encode())
12        response = b""
13        while True:
14            chunk = s.recv(4096)
15            if not chunk:
16                break
17            response += chunk
18        return response.decode()
19    except Exception as e:
20        return f"Error connecting to Tor Control Port: {e}"
21
22 def get_ports():
23     ports = {}
24     ports_file = "/tmp/nullexit-ports.env"
25     if os.path.exists(ports_file):
26         with open(ports_file, 'r') as f:
27             for line in f:
28                 if line.startswith("export "):
29                     # Split at first '=' to handle potential multiple '=' in value
30                     parts = line.replace("export ", "").strip().split("=", 1)
31                     if len(parts) == 2:
32                         key, val = parts
33                         try:
34                             ports[key] = int(val)
35                         except ValueError:
36                             ports[key] = val
37     return ports
38
39 def get_password(ports):
40     # 1. Try ports.env
41     password = ports.get("TOR_PASSWORD")
42     if password:
43         return password
44
45     # 2. Try .env
46     if os.path.exists(".env"):
47         try:
48             with open(".env", "r") as f:
49                 for line in f:
50                     if line.startswith("ADGUARD_PASSWORD="):
51                         return line.split("=", 1)[1].strip().strip('\\\'')
52                     if line.startswith("TOR_PASSWORD="):
53                         return line.split("=", 1)[1].strip().strip('\\\'')
54         except Exception:
55             pass
56
57     # 3. Fallback
58     return "nullexit"
59
60 def main():
61     ports = get_ports()
62     control_port = ports.get("TOR_CONTROL_PORT", 9051)
63     password = get_password(ports)
64
65     print(f"Connecting to Tor Control Port on 127.0.0.1:{control_port}...")
66     status = query_tor(control_port, "GETINFO circuit-status", password)
67
68     if "Error" in status:
69         print(status)
70         return
71
72     print("\nActive Tor Circuits:")
73     circuits = re.findall(r"(\d+)\s+BUILT\s+(.+) \s+PURPOSE", status)
74     if not circuits:
75         print("No active built circuits found yet. Tor might still be bootstrapping.")
76     return
77
```

```

78     for circ_id, path_str in circuits:
79         print(f"\nCircuit #{circ_id}:")
80         nodes = path_str.split(",")
81         for idx, node in enumerate(nodes):
82             fingerprint = node.split("~")[0].replace("$", "")
83             nickname = node.split("~")[1] if "~" in node else "unknown"
84
85             # Fetch node IP
86             ns_info = query_tor(control_port, f"GETINFO ns/id/{fingerprint}", password)
87             ip = "unknown"
88             for line in ns_info.split("\n"):
89                 if line.startswith("r "):
90                     parts = line.split(" ")
91                     if len(parts) >= 7:
92                         ip = parts[6]
93                         break
94
95             # Fetch country
96             country = "unknown"
97             if ip != "unknown":
98                 country_info = query_tor(control_port, f"GETINFO ip-to-country/{ip}", password)
99                 for line in country_info.split("\n"):
100                     if line.startswith("250-ip-to-country/"):
101                         country = line.split("=")[1].strip().upper()
102                         break
103
104             role = "Guard" if idx == 0 else ("Exit" if idx == len(nodes) - 1 else "Middle")
105             print(f" [{role}] {nickname} ({ip}) - Country: {country}")
106
107 if __name__ == "__main__":
108     main()

```

35 scripts/device-scramble.sh

```

1  #!/bin/bash
2  # scripts/device-scramble.sh  present a random, anonymous DEVICE identity on Wi-Fi.
3  #
4  # Randomizes the identifiers a network/observer uses to fingerprint your device:
5  #   hostname      ComputerName / LocalHostName / HostName (the DHCP hostname +
6  #                 the mDNS/Bonjour ".local" name broadcast on the LAN)
7  #   MAC (opt)    the layer-2 device address
8  #
9  # WHERE IT HELPS: open / no-auth Wi-Fi (caf, hotel, airport) where these IDs are
10 # the ONLY handle on you, and against co-located sniffers on any network. It does
11 # NOT hide you from a network you 802.1X-authenticate to that login ties every
12 # session to your account regardless of MAC (see the FaceTime/threat-model notes).
13 #
14 # SAFETY: 'scramble' saves your real identity first; 'restore' puts it back.
15 # 'test-mac' is FAIL-SAFE it always restores your original MAC at the end, and
16 # detects Apple-Silicon spoof-blocking and MAC-gated networks without stranding you.
17 #
18 # Usage (privileged actions need sudo):
19 #   bash scripts/device-scramble.sh status          # show current + saved identity (no sudo)
20 #   sudo bash scripts/device-scramble.sh scramble  # random hostname now (add --mac to also rotate MAC)
21 #   sudo bash scripts/device-scramble.sh restore   # put your real identity back
22 #   sudo bash scripts/device-scramble.sh test-mac  # can this network take a random MAC? (auto-reverts)
23 #   bash scripts/device-scramble.sh --dry-run      # print the random values it would use (no sudo, no change)
24
25 set -uo pipefail
26 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
27 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
28 cd "$PROJECT_ROOT" || exit 1
29 IFACE=en0
30 SAVE="$PROJECT_ROOT/.device-identity.orig" # gitignored
31 WARP_INHIBIT="/tmp/nullexit-warp-inhibit.marker" # WARP Watcher skips shutdown while this exists
32
33 rand_mac() { # locally-administered (0x02) + unicast (clear 0x01) first octet
34     printf '%02x:%02x:%02x:%02x:%02x:%02x' \
35         $(( (RANDOM & 0xFC) | 0x02 )) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM&255))
36 }
37 rand_host() { printf 'MacBook-%04X' $((RANDOM & 0xFFFF)); } # Apple-ish, blends in
38

```

```

39 cur_mac() { ifconfig "$IFACE" ether 2>/dev/null | awk '/ether/{print $2}'; }
40 cur_ip() { ipconfig getifaddr "$IFACE" 2>/dev/null; }
41 need_root() { [ "${id -u}" = 0 ] || { echo "device-scramble: this needs root  run with: sudo bash $0 $"
    >&&2; exit 1; }; }
42
43 cmd_status() {
44     echo "current:"
45     echo "  ComputerName  : $(scutil --get ComputerName 2>/dev/null)"
46     echo "  LocalHostName : $(scutil --get LocalHostName 2>/dev/null)"
47     echo "  HostName       : $(scutil --get HostName 2>/dev/null || echo '(unset)')"
48     echo "  MAC ($IFACE)   : $(cur_mac)   IP: $(cur_ip)"
49     if [ -f "$SAVE" ]; then echo; echo "saved original (restore target):"; sed 's/^/ /' "$SAVE"; fi
50 }
51
52 cmd_dry_run() {
53     echo "dry-run  would set (no changes made, no sudo needed):"
54     echo "  hostname  $(rand_host)"
55     echo "  MAC       $(rand_mac)   (locally-administered, unicast)"
56 }
57
58 save_original() {
59     [ -f "$SAVE" ] && return 0 # don't overwrite a real original with an already-scrambled one
60     {
61         echo "ComputerName=$(scutil --get ComputerName 2>/dev/null)"
62         echo "LocalHostName=$(scutil --get LocalHostName 2>/dev/null)"
63         echo "HostName=$(scutil --get HostName 2>/dev/null)"
64         echo "MAC=$(cur_mac)"
65     } > "$SAVE"
66 }
67
68 cmd_scramble() {
69     need_root scramble
70     save_original
71     local h; h=$(rand_host)
72     echo "device-scramble: hostname  $h"
73     scutil --set ComputerName "$h"
74     scutil --set LocalHostName "$h"
75     scutil --set HostName "$h"
76     dscacheutil -flushcache 2>/dev/null || true
77     if [ "${1:-}" = "--mac" ]; then
78         local m; m=$(rand_mac)
79         echo "device-scramble: MAC $m (disassociate set reconnect)"
80         networksetup -setairportpower "$IFACE" off; sleep 3 # can't set MAC while associated
81         ifconfig "$IFACE" ether "$m" 2>/dev/null
82         [ "${cur_mac}" = "$m" ] || echo " [!] MAC did not change  this driver refuses ifconfig spoofing on
            this hardware"
83         networksetup -setairportpower "$IFACE" on
84     fi
85     echo "done. Restore your real identity with: sudo bash $0 restore"
86 }
87
88 cmd_restore() {
89     need_root restore
90     [ -f "$SAVE" ] || { echo "device-scramble: no saved identity at $SAVE  nothing to restore."; exit 1; }
91     local cn lh hn mac
92     cn=$(sed -n 's/^ComputerName=//p' "$SAVE"); lh=$(sed -n 's/^LocalHostName=//p' "$SAVE")
93     hn=$(sed -n 's/^HostName=//p' "$SAVE"); mac=$(sed -n 's/^MAC=//p' "$SAVE")
94     [ -n "$cn" ] && scutil --set ComputerName "$cn"
95     [ -n "$lh" ] && scutil --set LocalHostName "$lh"
96     [ -n "$hn" ] && scutil --set HostName "$hn"
97     if [ -n "$mac" ] && [ "${cur_mac}" != "$mac" ]; then
98         networksetup -setairportpower "$IFACE" off; sleep 3 # disassociate before setting MAC
99         ifconfig "$IFACE" ether "$mac" 2>/dev/null
100        networksetup -setairportpower "$IFACE" on
101    fi
102    rm -f "$SAVE"
103    echo "restored: ComputerName=$cn  MAC=$(cur_mac)"
104 }
105
106 # The safe experiment: does THIS network accept a random MAC? Always reverts.
107 # CRITICAL: macOS refuses to change the MAC while the interface is ASSOCIATED to
108 # an SSID, so we disassociate (Wi-Fi radio off) BEFORE setting it, then reconnect.
109 # (This is the fix for the naive first attempt that set it while connected.)
110 cmd_test_mac() {
111     need_root test-mac
112     # Protect the running gateway: inhibit the WARP Watcher so the brief Wi-Fi blip
113     # isn't counted as tunnel death (which would auto-shutdown a healthy gateway).
114     # Only manage the marker if we created it  don't clobber recover.sh's.
115     local owned_inhibit=0
116     if [ ! -f "$WARP_INHIBIT" ]; then touch "$WARP_INHIBIT" && owned_inhibit=1; fi
117     trap '[ "$owned_inhibit" = 1 ] && rm -f "$WARP_INHIBIT"; EXIT INT TERM

```

```

118 echo "gateway guard: WARP Watcher inhibited during the test (containers stay up regardless)."
```

```

119 local orig; orig=$(cur_mac)
```

```

120 [ -n "$orig" ] || { echo "no $IFACE MAC found"; exit 1; }
```

```

121 echo "$orig" > /tmp/nullexit-mac.orig
```

```

122 local rnd; rnd=$(rand_mac)
```

```

123 echo "original MAC: $orig      trying random: $rnd"
```

```

124 echo "step 1: disassociate (Wi-Fi off)  you can't set the MAC while connected"
```

```

125 networksetup -setairportpower "$IFACE" off; sleep 3
```

```

126 echo "step 2: set MAC while disassociated"
```

```

127 ifconfig "$IFACE" ether "$rnd" 2>/dev/null
```

```

128 local after; after=$(cur_mac)
```

```

129 if [ "$after" != "$rnd" ]; then
```

```

130     echo "RESULT: MAC still $after even when disassociated  this driver refuses ifconfig spoofing on neo."
```

```

131     echo "          (Only Apple's native 'Private Wi-Fi Address' or a USB Wi-Fi adapter would work.)"
```

```

132     networksetup -setairportpower "$IFACE" on; exit 0
```

```

133 fi
```

```

134 echo "  MAC CHANGED to $after with disassociate-first. Reconnecting to test the network"
```

```

135 networksetup -setairportpower "$IFACE" on
```

```

136 local ip=""; for _ in $(seq 1 15); do sleep 2; ip=$(cur_ip); [ -n "$ip" ] && break; done
```

```

137 if [ -n "$ip" ]; then
```

```

138     echo "RESULT: SUCCESS  random MAC works AND got IP $ip. neo can spoof + this network accepts it."
```

```

139 else
```

```

140     echo "RESULT:  MAC spoof works, but no IP in ~30s  this network gates on your registered MAC (or
```

```

141       802.1X re-auth needed)."
```

```

142 fi
```

```

143 echo "restoring original MAC $orig "
```

```

144 networksetup -setairportpower "$IFACE" off; sleep 3
```

```

145 ifconfig "$IFACE" ether "$orig" 2>/dev/null
```

```

146 networksetup -setairportpower "$IFACE" on
```

```

147 ip=""; for _ in $(seq 1 15); do sleep 2; ip=$(cur_ip); [ -n "$ip" ] && break; done
```

```

148 echo "restored: MAC=$(cur_mac)  IP=${ip:-<reconnecting>}"
```

```

149 }
```

```

150 case "${1:-status}" in
```

```

151     status)    cmd_status ;;
```

```

152     --dry-run) cmd_dry_run ;;
```

```

153     scramble)  cmd_scramble "${2:-}" ;;
```

```

154     restore)   cmd_restore ;;
```

```

155     test-mac)  cmd_test_mac ;;
```

```

156     --help|-h) sed -n '2,26p' "$0" ;;
```

```

157     *) echo "Unknown: $1 (status|scramble [--mac]|restore|test-mac|--dry-run)" >&2; exit 2 ;;
```

```

158 esac
```

36 scripts/diagnose-host-leak.sh

```

1 #!/bin/bash
```

```

2 # scripts/diagnose-host-leak.sh  nullexit host-routing diagnostic
```

```

3 #
```

```

4 # Symptom this addresses: tests like https://www.whatismyip.com on the
```

```

5 # *HOST MAC* return the underlying physical ISP/ASN
```

```

6 # local ISP / campus network") even though the nullexit gateway is active and remote
```

```

7 # tailnet clients correctly egress via Cloudflare WARP. The container and
```

```

8 # the gateway itself are healthy  only the host's own routing is leaking.
```

```

9 #
```

```

10 # This script runs 7 checks in sequence, classifies the result into ONE
```

```

11 # of the three known leak scenarios, writes the full report to a
```

```

12 # timestamped file in the project root, and prints a verdict + a
```

```

13 # ready-to-run fix command on stdout. Pass --fix to apply the matching
```

```

14 # remediation in the same invocation and re-verify egress afterwards.
```

```

15 # See devref.md 10.30 for the deep-dive rationale.
```

```

16 #
```

```

17 # Usage:
```

```

18 #   bash scripts/diagnose-host-leak.sh           # diagnose + write report
```

```

19 #   bash scripts/diagnose-host-leak.sh --fix      # diagnose + fix + re-verify
```

```

20 #   bash scripts/diagnose-host-leak.sh --watch    # full baseline then loop checks every 60s
```

```

21 #   bash scripts/diagnose-host-leak.sh --watch 30 # same, with custom interval (seconds)
```

```

22 #   bash scripts/diagnose-host-leak.sh --help     # this message
```

```

23 #
```

```

24 # Multiple runs produce timestamped files (one per run) so historical
```

```

25 # reports can be diffed. The newest file is the most recent.
```

```

26
```

```

27 set -uo pipefail
```

```

28 # NOTE: errexit ('set -e') is intentionally NOT enabled. Several checks
```

```

29 # below are EXPECTED to fail and that's diagnostic data (e.g. 'tailscale
```

```

30 # status' returns non-zero when the daemon is wedged). Every command
```

```

31 # pipes to a helper that records the exit code without killing the script.
32
33 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
34 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
35 TIMESTAMP="$(date -u +%Y%m%dT%H%M%S)"
36 source "$SCRIPT_DIR/common.sh"
37 LOG_FILE="$PROJECT_ROOT/output.log"
38 REPORT_FILE="$PROJECT_ROOT/host-leak-diagnostic-$TIMESTAMP.txt"
39
40 # Argument parsing
41 DO_FIX=false
42 DO_WATCH=false
43 WATCH_INTERVAL=60
44 WATCH_LOG="$PROJECT_ROOT/host-leak-watch.log"
45 case "${1:-}" in
46   --fix) DO_FIX=true ;;
47   --watch)
48     DO_WATCH=true
49     # Optional second argument: custom interval in seconds
50     if [ -n "${2:-}" ] && [[ "$2" =~ ^[0-9]+$ ]]; then
51       WATCH_INTERVAL="$2"
52     fi
53     ;;
54   --help|-h)
55     sed -n '2,31p' "$0"
56     exit 0
57     ;;
58   *)
59     : # default: diagnose only
60     ;;
61   *)
62     echo "Unknown argument: $1" >&2
63     echo "Run with --help for usage." >&2
64     exit 2
65     ;;
66 esac
67
68 # Colours (auto-disabled when stdout is not a tty, e.g. piped)
69 if [ -t 1 ]; then
70   RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
71   BOLD='\033[1m'; NC='\033[0m'
72 else
73   RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
74 fi
75
76 # Output helpers write to BOTH stdout and the report file
77 exec 3>> "$REPORT_FILE" || {
78   printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
79   exit 1
80 }
81 section() {
82   local title="$1"
83   printf '\n%b %s %b\n' "$BOLD" "$title" "$NC" >&3
84   printf '\n%b %s %b\n' "$BOLD" "$title" "$NC"
85 }
86 line() { printf ' %b\n' "$1" >&3; printf ' %b\n' "$1"; }
87 ok() { line "${GREEN}${NC} $1"; }
88 warn() { line "${YELLOW}${NC} $1"; }
89 fail() { line "${RED}${NC} $1"; }
90 note() { line "(no $1)"; }
91
92 # Watch-mode header (only shown when entering watch loop)
93 if [ "$DO_WATCH" = "true" ]; then
94   echo ""
95   echo ""
96   echo "n u l l e x i t   w a t c h   m o d e"
97   echo ""
98   echo "Interval:      ${WATCH_INTERVAL}s"
99   echo "Watch log:     $WATCH_LOG"
100   echo ""
101 else
102   echo ""
103   echo "n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c"
104   echo ""
105   echo "Timestamp (UTC): $TIMESTAMP"
106   echo "Project root:    $PROJECT_ROOT"
107   echo "Report file:     $REPORT_FILE"
108   echo "Mode:           $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only (
109     pass --fix to remediate)")"
110   echo ""
111   echo "" >&3

```



```

111 echo "" >&3
112 echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c" >&3
113 echo "" >&3
114 echo " Timestamp (UTC): $TIMESTAMP" >&3
115 echo " Mode:          $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only")"
    >&3
116 fi
117
118 # Always add Homebrew to PATH (same as toggle.sh / recover.sh)
119 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
120
121 # Pure-bash timeout (macOS lacks GNU 'timeout')
122 run_with_timeout() {
123     local timeout_sec="$1"
124     shift
125     if [ "$#" -eq 0 ]; then return 1; fi
126     "$@" &
127     local cmd_pid=$!
128     (
129         sleep "$timeout_sec"
130         if kill -0 "$cmd_pid" 2>> "$LOG_FILE"; then
131             kill -9 "$cmd_pid" 2>> "$LOG_FILE" || true
132         fi
133     ) >> "$LOG_FILE" 2>&1 &
134     local watcher_pid=$!
135     set +e
136     wait "$cmd_pid"
137     local exit_status=$?
138     set -uo pipefail
139     kill "$watcher_pid" 2>> "$LOG_FILE" || true
140     return "$exit_status"
141 }
142
143 #
144 # Quick-check functions (used by --watch loop; return simple status strings)
145 # Each echoes its result so the caller can capture and compare.
146 #
147
148 quick_warp() {
149     # Returns: "on", "off", "unknown"
150     local out
151     out=$(run_with_timeout 8 curl -s --max-time 6 \
152         https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null || echo "")
153     local warp
154     warp=$(printf '%s' "$out" | awk -F'= ' '/^warp={print $2; exit}')
155     printf '%s' "${warp:-unknown}"
156 }
157
158 quick_ipv6() {
159     # Returns: the IPv6 address if leaking, or empty string if clean
160     local out
161     out=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>/dev/null || echo "")
162     if [ -n "$out" ] && printf '%s' "$out" | grep -qE '^[0-9a-fA-F:]+$'; then
163         printf '%s' "$out"
164     fi
165 }
166
167 quick_default_route() {
168     # Returns: "utun" if default route is via Tailscale utun*,
169     #           "wifi" if via physical Wi-Fi (192.168.x.x),
170     #           "other" otherwise
171     local route
172     route=$(netstat -rn 2>/dev/null | awk '/^default/ {print $0; exit}')
173     if [ -z "$route" ]; then
174         printf '%s' "none"
175     elif printf '%s' "$route" | grep -qE 'utun[0-9]+'; then
176         printf '%s' "utun"
177     elif printf '%s' "$route" | grep -qE '^(default|0\.0\.0\.0).*\b192\.168\.'; then
178         printf '%s' "wifi"
179     else
180         printf '%s' "other"
181     fi
182 }
183
184 # Watch loop (runs after baseline diagnostic when --watch is passed)
185 run_watch_loop() {
186     local baseline_warp="$1"
187     local baseline_default="$2"
188     local baseline_ipv6_leak="$3"
189     local interval="$4"
190

```

```

191 echo ""
192 echo ""
193 echo " Baseline established. Monitoring every ${interval}s..."
194 echo " warp:      $baseline_warp"
195 echo " default:    $baseline_default"
196 echo ""
197
198 # Safety: warn if interval is very short (avoids hammering APIs)
199 if [ "$interval" -lt 30 ]; then
200     printf '%b Short interval (%ss) rate-limiting risk on external APIs\b\n' "$YELLOW" "$interval" "$NC"
201     printf ' Consider 30s or longer for production use.\n'
202 fi
203 printf '[%s] WATCH START warp=%s default=%s interval=%ss\n' \
204 "$date -u +%Y-%m-%dT%H:%M:%SZ)" "$baseline_warp" "$baseline_default" "$interval" \
205 >> "$WATCH_LOG"
206
207 local iteration=0
208 local last_warp="$baseline_warp"
209 local last_default="$baseline_default"
210 local last_ipv6_ok=$(( [ "$baseline_ipv6_leak" = "false" ] && echo true || echo false))
211
212 trap 'printf "\n\bWatch stopped by user. Log at: %s\b\n" "$YELLOW" "$WATCH_LOG" "$NC"; exit 0' INT TERM
213
214 while true; do
215     iteration=$((iteration + 1))
216
217     local warp_now ipv6_now default_now
218     warp_now=$(quick_warp)
219     ipv6_now=$(quick_ipv6)
220     default_now=$(quick_default_route)
221
222     local ts
223     ts=$(date -u +%H:%M:%S)
224
225     # WARP check
226     if [ "$warp_now" != "$last_warp" ]; then
227         if [ "$warp_now" = "off" ] || [ "$warp_now" = "unknown" ]; then
228             printf '\n\b ALERT %b\n' "$RED$BOLD" "$NC"
229             printf '%b[%s] %bWARP FLIPPED: %s %s\b YOUR IP IS LEAKING!\n' \
230 "$RED" "$ts" "$BOLD" "$last_warp" "$warp_now" "$NC"
231             printf '[%s] ALERT warp flipped %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
232         else
233             printf '\n\b[%s] warp recovered: %s %s\b\n' \
234 "$GREEN" "$ts" "$last_warp" "$warp_now" "$NC"
235             printf '[%s] OK warp recovered %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
236         fi
237         last_warp="$warp_now"
238     fi
239
240     # IPv6 check
241     if [ -n "$ipv6_now" ]; then
242         if [ "$last_ipv6_ok" = true ]; then
243             printf '\n\b ALERT %b\n' "$RED$BOLD" "$NC"
244             printf '%b[%s] %bIPv6 LEAK DETECTED %s\b\n' \
245 "$RED" "$ts" "$BOLD" "$ipv6_now" "$NC"
246             printf '[%s] ALERT IPv6 leak %s\n' "$ts" "$ipv6_now" >> "$WATCH_LOG"
247             last_ipv6_ok=false
248         fi
249     else
250         if [ "$last_ipv6_ok" = false ]; then
251             printf '%b[%s] IPv6 leak resolved\b\n' "$GREEN" "$ts" "$NC"
252             printf '[%s] OK IPv6 resolved\n' "$ts" >> "$WATCH_LOG"
253         fi
254         last_ipv6_ok=true
255     fi
256
257     # Default route check
258     if [ "$default_now" != "$last_default" ]; then
259         if [ "$default_now" != "utun" ]; then
260             printf '\n\b ALERT %b\n' "$RED$BOLD" "$NC"
261             printf '%b[%s] %bDEFAULT ROUTE CHANGED: %s %s\b traffic may bypass WARP!\n' \
262 "$RED" "$ts" "$BOLD" "$last_default" "$default_now" "$NC"
263             printf '[%s] ALERT route changed %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
264         else
265             printf '%b[%s] default route recovered: %s %s\b\n' \
266 "$GREEN" "$ts" "$last_default" "$default_now" "$NC"
267             printf '[%s] OK route recovered %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
268         fi
269         last_default="$default_now"
270     fi

```

```

271
272 # Heartbeat (every 10th iteration or if all OK)
273 if [ $((iteration % 10)) -eq 0 ]; then
274     printf "[%s] #d warp=%s ipv6=%s route=%s\n" \
275         "$ts" "$iteration" "$warp_now" "${ipv6_now:-none}" "$default_now" >> "$WATCH_LOG"
276     printf '\r%b[%s] #d warp=%s route=%s (Ctrl+C to stop)%b' \
277         "$GREEN" "$ts" "$iteration" "$warp_now" "$default_now" "$NC"
278 fi
279
280     sleep "$interval"
281 done
282 }
283
284 #
285 # Begin main diagnostic (skipped in non-watch modes if we're just watching)
286 #
287
288 ACTIVE_SERVICE=$(get_active_service)
289 ENO_SERVICE=$(networksetup -listnetworkserviceorder 2>> "$LOG_FILE" \
290     | grep -B 1 "Device: en0" | head -1 \
291     | sed -E 's/^\([0-9*\]\)\s*//' || true)
292 [ -z "$ENO_SERVICE" ] && ENO_SERVICE="Wi-Fi"
293
294 # Resolve the gateway's Tailscale IP.
295 GATEWAY_TS_IP=""
296 if [ -f "$PROJECT_ROOT/.gateway_ip" ]; then
297     GATEWAY_TS_IP=$(cat "$PROJECT_ROOT/.gateway_ip" 2>> "$LOG_FILE" \
298         | tr -d '\r' | awk 'NR==1{print $1; exit}')
299 fi
300 if [ -z "$GATEWAY_TS_IP" ] && command -v docker >> "$LOG_FILE" 2>&1 \
301     && docker compose ps --status running 2>/dev/null | grep -q 'tailscale'; then
302     GATEWAY_TS_IP=$(run_with_timeout 10 docker compose exec -T tailscale \
303         tailscale ip -4 2>> "$LOG_FILE" \
304         | tr -d '\r' | awk 'NR==1{print $1; exit}')
305 fi
306
307 # Scenario classification state
308 SCENARIO=""
309 WARP_STATUS=""
310 SOCKS5_ENABLED=false
311 TS_ON_MESH=false
312 DEFAULT_VIA_UTUN=false
313 IPV6_LEAK=false
314
315 #
316 section "1/8 Host Tailscale mesh status"
317 #
318 if ! command -v tailscale >> "$LOG_FILE" 2>&1; then
319     fail "tailscale CLI not on PATH install via 'brew install tailscale' then re-run"
320     echo " (a Tailscale daemon is required for the exit-node path to work)"
321 else
322     set +e
323     TS_OUT=$(run_with_timeout 5 tailscale status 2>&1)
324     TS_EXIT=$?
325     set -uo pipefail
326     TS_HEAD=$(printf '%s\n' "$TS_OUT" | head -5 | tr -d '\r')
327     if [ "$TS_EXIT" -eq 137 ]; then
328         fail "tailscaled daemon is wedged/unresponsive (5s timeout)"
329         echo " Fix path: 'sudo brew services restart tailscale'"
330     elif printf '%s' "$TS_OUT" | grep -qE 'Logged out|not logged in|expired'; then
331         fail "tailscale is logged out / auth expired on this host"
332         echo " Fix path: 'sudo tailscale up' (browser auth once)"
333     elif printf '%s' "$TS_OUT" | grep -q '100\.'; then
334         ok "Host is on tailnet (100.x.x.x visible in status)"
335         TS_ON_MESH=true
336         if printf '%s' "$TS_OUT" | grep -q "$GATEWAY_TS_IP"; then
337             ok "Gateway $GATEWAY_TS_IP visible in host's peer list"
338         elif [ -n "$GATEWAY_TS_IP" ]; then
339             warn "Gateway $GATEWAY_TS_IP NOT in host's peer list possible auth/key mismatch"
340         fi
341         line " first 5 lines of 'tailscale status' "
342         line "$(${printf '%s' "$TS_HEAD" | awk '{print " " $0}'}"
343     else
344         fail "tailscale output unrecognized:"
345         line "$(${printf '%s' "$TS_OUT" | head -3 | awk '{print " " $0}'}"
346     fi
347 fi
348
349 #
350 section "2/8 Active SOCKS5 proxy on Wi-Fi"
351 #

```

```

352 SOCKS_OUT=$(networksetup -getsocksfirewallproxy "$ACTIVE_SERVICE" 2>> "$LOG_FILE" \
353 || networksetup -getsocksfirewallproxy Wi-Fi 2>> "$LOG_FILE" \
354 || echo "Could not query SOCKS5 proxy state")
355 if printf '%s' "$SOCKS_OUT" | grep -q "Enabled: Yes"; then
356     fail "SOCKS5 proxy IS enabled on $ACTIVE_SERVICE toggle.sh fell back to SOCKS5 mode last run"
357     SOCKS5_ENABLED=true
358     line " Full state:"
359     line "$(printf '%s' "$SOCKS_OUT" | awk '{print "      $0}')"
360     line " Browse on macOS ignores system SOCKS5 by default traffic bypasses WARP."
361 else
362     ok "SOCKS5 proxy is OFF on $ACTIVE_SERVICE exit-node path should be active"
363 fi
364
365 #
366 section "3/8 Egress IP + WARP tunnel verification (definitive test)"
367 #
368 CDN_OUT=$(run_with_timeout 8 curl -s --max-time 6 \
369 https://www.cloudflare.com/cdn-cgi/trace 2>&1 || echo "(curl failed)")
370 CDN_WARP=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^warp={print $2; exit}')
371 CDN_IP=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^ip={print $2; exit}')
372 CDN_COLO=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^colo={print $2; exit}')
373 WARP_STATUS="$CDN_WARP"
374 if [ "$CDN_WARP" = "on" ]; then
375     ok "warp=on tunnel is hot, host traffic IS going through Cloudflare WARP"
376     line " Public IP: $CDN_IP"
377     line " Cloudflare: ${CDN_COLO:-unknown colo}"
378     line " whatismyip.com's 'local ISP' result is its ISP-database error, not yours."
379 elif [ "$CDN_WARP" = "off" ]; then
380     fail "warp=off host traffic is NOT going through WARP"
381     line " Public IP: $CDN_IP (NOT Cloudflare)"
382     warn "This is the actual leak."
383 else
384     warn "could not determine warp status (curl failed or returned unexpected shape)"
385     line " Raw output: $(printf '%s' "$CDN_OUT" | head -3 | tr '\n' '|')"
386 fi
387
388 #
389 section "4/8 IPv6 leak probe (bypasses IPv4 exit-node on campus networks)"
390 #
391 IPV6_OUT=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
392 if [ -n "$IPV6_OUT" ] && printf '%s' "$IPV6_OUT" | grep -qE '[0-9a-fA-F:]+$'; then
393     fail "host has live IPv6 egress $IPV6_OUT (A6 record bypasses WARP)"
394     IPV6_LEAK=true
395     line " Tailscale exit-node advertisement is IPv4-only, so IPv6 traffic skips it."
396     line " Quick-fix: 'sudo networksetup -setv6off $ACTIVE_SERVICE'"
397 else
398     ok "no IPv6 egress detected (good IPv6 won't bypass the tunnel)"
399 fi
400
401 #
402 section "5/8 Host egress route for real traffic (route -n get 1.1.1.1)"
403 #
404 # WHY NOT 'netstat -rn | grep default':
405 # On macOS, Tailscale does NOT replace the physical default route installed by
406 # DHCP. Instead it adds a second interface-scoped route bound to the utun
407 # interface that the kernel prefers for real traffic forwarding. The literal
408 # "default" entry (192.168.x.x via en0) is left untouched in the table as a
409 # BSD routing quirk it's not lying, it's just answering a different question.
410 # The correct question is: "where does a real destination actually resolve?"
411 # 'route -n get 1.1.1.1' asks the kernel's forwarding logic directly.
412 ROUTE_GET=$(route -n get 1.1.1.1 2>/dev/null)
413 ROUTE_IFACE=$(echo "$ROUTE_GET" | awk '/interface:{print $2}')
414 ROUTE_GATEWAY=$(echo "$ROUTE_GET" | awk '/gateway:{print $2}')
415 if [ -z "$ROUTE_IFACE" ]; then
416     warn "route -n get 1.1.1.1 returned no interface routing table may be empty"
417 else
418     line " interface: $ROUTE_IFACE gateway: ${ROUTE_GATEWAY:-n/a}"
419     if echo "$ROUTE_IFACE" | grep -qE '^utun[0-9]+'; then
420         ok "1.1.1.1 resolves via $ROUTE_IFACE Tailscale interface-scoped route is winning"
421         DEFAULT_VIA_UTUN=true
422     elif echo "$ROUTE_IFACE" | grep -qE '^en[0-9]+'; then
423         fail "1.1.1.1 resolves via $ROUTE_IFACE (physical Wi-Fi) exit-node interface route not active"
424     else
425         warn "1.1.1.1 resolves via unexpected interface: $ROUTE_IFACE"
426     fi
427 fi
428
429 #
430 section "6/8 Last toggle.sh START-relevant lines in output.log (if any)"
431 #
432 if [ ! -f "$LOG_FILE" ]; then

```

```

433 note "output.log"
434 line " toggle.sh was never run from this directory. Run './toggle.sh' first,"
435 line " then re-run this diagnostic."
436 else
437 HITS=$(grep -E 'Exit node enabled|SKIP_EXIT_NODE|Pre-flight|Falling back to SOCKS|Starting|Successfully
    ' \
438     "$LOG_FILE" 2>/dev/null | tail -8 || true)
439 if [ -z "$HITS" ]; then
440     note "matching lines in output.log"
441 else
442     line "$(printf '%s\n' "$HITS" | awk '{print "    "$0}')"
443 fi
444 fi
445
446 #
447 section "7/8 Gateway IP we should be pointing at"
448 #
449 if [ -n "$GATEWAY_TS_IP" ]; then
450     ok "gateway Tailscale IP: $GATEWAY_TS_IP"
451 else
452     warn "could not resolve gateway Tailscale IP"
453     line " (looked in .gateway_ip and tried docker compose exec tailscale ip -4)"
454 fi
455
456 #
457 section "8/8 Tailscale System Extension (App Store conflict check)"
458 #
459 SE_PENDING_UNINSTALL=false
460 SE_ACTIVE=false
461 if command -v systemextensionsctl >/dev/null 2>&1; then
462     SE_OUT=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale')
463     if [ -n "$SE_OUT" ]; then
464         line " Tailscale System Extension(s) detected:"
465         line "$(printf '%s' "$SE_OUT" | awk '{print "    "$0}')"
466         if printf '%s' "$SE_OUT" | grep -q 'waiting to uninstall'; then
467             SE_PENDING_UNINSTALL=true
468             warn "System Extension is pending uninstall on next reboot."
469         else
470             SE_ACTIVE=true
471             warn "Active Tailscale System Extension detected. This will conflict with Homebrew Tailscale."
472         fi
473     else
474         ok "no Tailscale System Extension detected"
475     fi
476 else
477     note "systemextensionsctl not available (not on macOS or permission restricted)"
478 fi
479
480 #
481 section "CLASSIFICATION applies ONE of the four known scenarios"
482 #
483
484 if [ "$WARP_STATUS" = "on" ] && [ "$IPV6_LEAK" = "false" ]; then
485     SCENARIO="OK"
486 elif [ "$SE_PENDING_UNINSTALL" = "true" ]; then
487     SCENARIO="SE_PENDING"
488 elif [ "$SOCKS5_ENABLED" = "true" ]; then
489     SCENARIO="A"
490 elif [ "$IPV6_LEAK" = "true" ]; then
491     SCENARIO="B"
492 elif [ "$DEFAULT_VIA_UTUN" = "false" ] && [ "$TS_ON_MESH" = "true" ]; then
493     SCENARIO="C"
494 else
495     SCENARIO="UNKNOWN"
496 fi
497
498 case "$SCENARIO" in
499     SE_PENDING)
500         echo ""
501         echo -e " ${RED}${BOLD}VERDICT: Scenario SE_PENDING Tailscale System Extension Pending Uninstall${NC}"
502         echo " A prior App Store Tailscale System Extension is pending uninstall."
503         echo " Since System Integrity Protection (SIP) is enabled, macOS blocks manual"
504         echo " uninstallation of terminated system extensions until the next reboot."
505         echo " Brew-managed tailscaled cannot engage the Network Extension slot until this is resolved."
506         echo ""
507         echo " ${BOLD}Recommended fix:${NC}"
508         echo " sudo shutdown -r now"
509         echo " # After reboot, run: ./toggle.sh"
510     ;;
511     A)

```

```

512 echo ""
513 echo -e " ${RED}${BOLD}VERDICT: Scenario A SOCKS5 fallback lane${NC}"
514 echo " toggle.sh's pre-flight checks failed last run. The script activated"
515 echo " SOCKS5 + local DNS proxy as a fallback. DNS gets hijacked but"
516 echo " browsers on macOS ignore the system SOCKS5 traffic egresses"
517 echo " direct through the local network."
518 echo ""
519 echo " ${BOLD}Recommended fix:${NC}"
520 echo "     echo 'GATEWAY_BYPASS_PING=true' >> $PROJECT_ROOT/.env"
521 echo "     sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
522 if [ -n "$GATEWAY_TS_IP" ]; then
523     echo "         --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
524 else
525     echo "         --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true # fill in $GATEWAY_TS_IP"
526 fi
527 echo "     sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
528 echo "     ./toggle.sh"
529 ;;
530 B)
531 echo ""
532 echo -e " ${RED}${BOLD}VERDICT: Scenario B IPv6 leak over dual-stack campus/local IPv6${NC}"
533 echo " Tailscale exit-node advertises only IPv4. macOS sends AAAA queries"
534 echo " straight out en0, which is dual-stack on most campus APs IPv6"
535 echo " traffic bypasses the WARP tunnel entirely."
536 echo ""
537 echo " ${BOLD}Recommended fix:${NC}"
538 echo "     sudo networksetup -setv6off $ACTIVE_SERVICE"
539 echo "     sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
540 echo "     # Re-enable IPv6 later with: sudo networksetup -setv6automatic $ACTIVE_SERVICE"
541 ;;
542 C)
543 echo ""
544 echo -e " ${RED}${BOLD}VERDICT: Scenario C Tailscale route-freeze${NC}"
545 echo " Host's default route points at the Wi-Fi gateway instead of the"
546 echo " Tailscale utun* interface. The exit-node preference is set but the"
547 echo " routing-table assertion did not take effect (devref 10.26)."

```

```

591 C L A S S I F I C A T I O N   R E S U L T
592
593 Scenario: $$SCENARIO
594 WARP egress:  $WARP_STATUS (cdn-cgi/trace)
595 SOCKS5 lane:  $([ "$SOCKS5_ENABLED" = true ] && echo "ENABLED (browsers bypass)" || echo "off")
596 IPv6 leak:    $([ "$IPV6_LEAK" = true ] && echo "YES (campus IPv6 egress)" || echo "no")
597 Default via:  $([ "$DEFAULT_VIA_UTUN" = true ] && echo "utun* (Tailscale)" || echo "Wi-Fi gateway")
598 Host on mesh: $([ "$TS_ON_MESH" = true ] && echo "yes" || echo "no")
599 SE pending:   $([ "$SE_PENDING_UNINSTALL" = true ] && echo "yes" || echo "no")
600 Gateway IP:   ${GATEWAY_TS_IP:-unresolved}
601 EOF
602 )
603 printf '%s\n' "$SECOND_BLOCK" >&3
604
605 #
606 # --fix mode
607 #
608 if [ "$DO_FIX" = "true" ] && [ "$SCENARIO" != "OK" ] && [ "$SCENARIO" != "UNKNOWN" ]; then
609     echo ""
610     echo ""
611     echo " A P P L Y I N G   F I X (scenario $SCENARIO)"
612     echo ""
613     echo ""
614
615     case "$SCENARIO" in
616         SE_PENDING)
617             warn "System Extension uninstall is pending on next reboot."
618             line "Please run 'sudo shutdown -r now' to reboot the system."
619             ;;
620         A)
621             if [ -n "$GATEWAY_TS_IP" ]; then
622                 line " writing GATEWAY_BYPASS_PING=true to .env"
623                 ENV="$PROJECT_ROOT/.env"
624                 [ -f "$ENV" ] || touch "$ENV"
625                 if grep -q 'GATEWAY_BYPASS_PING=' "$ENV" 2>/dev/null; then
626                     sed -i 's/^GATEWAY_BYPASS_PING=.*GATEWAY_BYPASS_PING=true/' "$ENV"
627                 else
628                     printf '\nGATEWAY_BYPASS_PING=true\n' >> "$ENV"
629                 fi
630                 line " re-asserting tailscale exit-node (with --reset)"
631                 sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
632                     --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
633                     >> "$LOG_FILE" 2>&1 \
634                     && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
635                     || warn "tailscale up did not return success (check $LOG_FILE)"
636             else
637                 warn "no gateway Tailscale IP resolved skipping tailscale up"
638                 warn "fix manually: 'sudo tailscale up --reset ... --exit-node=<TS_IP>'"
639             fi
640             ;;
641         B)
642             line " disabling IPv6 on $ACTIVE_SERVICE while gateway is up"
643             sudo -n networksetup -setv6off "$ACTIVE_SERVICE" >> "$LOG_FILE" 2>&1 \
644                 && ok "IPv6 disabled on $ACTIVE_SERVICE (re-enable with 'setv6automatic')" \
645                 || warn "networksetup -setv6off failed (check $LOG_FILE)"
646             ;;
647         C)
648             line " restarting tailscaled (route-table fix)"
649             run_with_timeout 15 brew services restart tailscale >> "$LOG_FILE" 2>&1 \
650                 || run_with_timeout 15 sudo -n brew services restart tailscale >> "$LOG_FILE" 2>&1
651             sleep 3
652             line " flushing stale host routes + DNS cache"
653             sudo -n route delete -net 0.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
654             sudo -n route delete -net 128.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
655             sudo -n dscacheutil -flushcache >> "$LOG_FILE" 2>&1 || true
656             sudo -n killall -HUP mDNSResponder >> "$LOG_FILE" 2>&1 || true
657             if [ -n "$GATEWAY_TS_IP" ]; then
658                 line " re-asserting exit-node after restart"
659                 sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
660                     --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
661                     >> "$LOG_FILE" 2>&1 \
662                     && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
663                     || warn "tailscale up did not return success (check $LOG_FILE)"
664             fi
665             ;;
666     esac
667
668     echo ""
669     echo "Re-verifying after fix..."
670     sleep 3
671     CDN_WARP_AFTER=$(run_with_timeout 8 curl -s --max-time 6 \

```



```

672     https://www.cloudflare.com/cdn-cgi/trace 2>&1 \
673     | awk -F=' ' '/^warp=/{print $2; exit}'')
674     IPV6_AFTER=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
675     printf '\n[AFTER FIX]\n' >&3
676     printf ' warp status:      %s\n' "${CDN_WARP_AFTER}" >&3
677     printf ' IPv6 egress:      %s\n' "${IPV6_AFTER:-none}" >&3
678     line " warp status:      ${CDN_WARP_AFTER:-unknown}"
679     line " IPv6 egress:      ${IPV6_AFTER:-none}"
680     if [ "${CDN_WARP_AFTER}" = "on" ] && [ -z "${IPV6_AFTER}" ]; then
681         echo ""
682         ok "Fix verified host is now routing through WARP. Remote clients unaffected."
683     else
684         echo ""
685         warn "Fix did not fully take effect on first try. Likely causes:"
686         line " tailscaled needed a longer settling window (re-run diagnostic in 30s)"
687         line " The classified scenario was wrong; paste this report for triage."
688         line " Tailscale needs a full auth refresh: 'sudo tailscale up'"
689     fi
690 fi
691
692 echo ""
693 echo ""
694 echo " Full report written to:"
695 echo " $REPORT_FILE"
696 echo ""
697 echo ""
698
699 #
700 # --watch mode: enter the continuous monitoring loop after baseline diagnostic
701 #
702 if [ "$DO_WATCH" = "true" ]; then
703     # Build baseline from the full diagnostic just completed
704     BASELINE_WARP="$WARP_STATUS"
705     BASELINE_DEFAULT=$( [ "${DEFAULT_VIA_UTUN}" = "true" ] && echo "utun" || echo "wifi/other")
706
707     if [ "$SCENARIO" != "OK" ]; then
708         warn "Baseline diagnostic found issues (scenario: $SCENARIO)."
709         line " Watch loop will still run alerts fire on any STATE CHANGE from current baseline."
710         line " Fix the issue first? Re-run with: bash scripts/diagnose-host-leak.sh --fix"
711         echo ""
712     fi
713
714     run_watch_loop "$BASELINE_WARP" "$BASELINE_DEFAULT" "$IPV6_LEAK" "$WATCH_INTERVAL"
715 fi
716
717 exit 0

```

37 scripts/dns-proxy.py

```

1  #!/usr/bin/env python3
2  """
3  dns-proxy.py Local DNS proxy for nullexit gateway.
4
5  Listens on UDP port 53, receives DNS queries, forwards them over TCP
6  to AdGuard at 127.0.0.1:5354 (the Docker port mapping), and returns
7  the response via UDP. Handles the DNS-over-TCP wire format (2-byte
8  length prefix) correctly unlike socat which just forwards raw bytes.
9  """
10
11 import socket
12 import struct
13 import sys
14 import os
15 from concurrent.futures import ThreadPoolExecutor
16
17 LISTEN_ADDR = os.environ.get("LISTEN_ADDR", "127.0.0.1")
18 LISTEN_PORT = int(os.environ.get("LISTEN_PORT", 53))
19 TARGET_HOST = os.environ.get("TARGET_HOST", "127.0.0.1")
20 TARGET_PORT = int(os.environ.get("TARGET_PORT", 5354))
21
22
23 def handle_query(data: bytes, client_addr: tuple, sock: socket.socket) -> None:
24     """Forward a single DNS query over TCP and send the response back via UDP."""
25     try:
26         tcp = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
27         tcp.settimeout(5)

```

```

28     tcp.connect((TARGET_HOST, TARGET_PORT))
29
30     # DNS over TCP: prepend 2-byte length (big-endian)
31     tcp.sendall(struct.pack("!H", len(data)) + data)
32
33     # Read the 2-byte response length
34     raw_len = tcp.recv(2)
35     if len(raw_len) < 2:
36         tcp.close()
37         return
38     resp_len = struct.unpack("!H", raw_len)[0]
39
40     # Read the full response
41     response = b""
42     while len(response) < resp_len:
43         chunk = tcp.recv(resp_len - len(response))
44         if not chunk:
45             break
46         response += chunk
47
48     tcp.close()
49
50     # Send the response back over UDP (strip TCP length prefix)
51     if response:
52         sock.sendto(response, client_addr)
53 except Exception:
54     # Silently drop failed queries (malformed, timeout, etc.)
55     pass
56
57
58 def main() -> None:
59     # Create UDP socket
60     sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
61     sock.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
62
63     try:
64         sock.bind((LISTEN_ADDR, LISTEN_PORT))
65     except PermissionError:
66         print(f"dns-proxy: ERROR need root to bind port {LISTEN_PORT}", flush=True)
67         sys.exit(1)
68     except OSError as e:
69         print(f"dns-proxy: ERROR {e}", flush=True)
70         sys.exit(1)
71
72     print(f"dns-proxy: listening on UDP {LISTEN_ADDR}:{LISTEN_PORT} TCP {TARGET_HOST}:{TARGET_PORT}",
73         flush=True)
74
75     # Use a thread pool to handle concurrent DNS queries efficiently without thread exhaustion
76     executor = ThreadPoolExecutor(max_workers=64)
77
78     while True:
79         try:
80             data, client_addr = sock.recvfrom(4096) # Support EDNS0 (up to 4096 bytes) to prevent
81             truncation
82             executor.submit(handle_query, data, client_addr, sock)
83         except KeyboardInterrupt:
84             break
85         except Exception:
86             pass
87
88     executor.shutdown(wait=False)
89     sock.close()
90
91 if __name__ == "__main__":
92     main()

```

38 scripts/dns_anomaly_detector.py

```

1 #!/usr/bin/env python3
2 """
3 nullexit DNS anomaly detector (v2)  report-only DNS-tunneling / DGA / C2 finder.
4
5 Replaces the v1 whole-FQDN Shannon-entropy + mean3 demo (which measured the
6 wrong scope, was fooled by short encoded labels, assumed a normal distribution
7 DNS names don't have, and ignored the volumetric signal that actually defines

```

```

8 tunneling). v2 keeps the "statistical, data-based, 100% local, no third-party
9 deps" spirit but detects far better:
10
11 1. n-gram (character bigram) improbability of the SUBDOMAIN labels only,
12    trained on YOUR querylog. Robust on short strings where Shannon entropy
13    fails (entropy is capped at log2(len)); scores "doesn't look like real
14    naming", which is what encoded tunneling labels are.
15 2. per-registered-domain AGGREGATION (unique-subdomain count, query volume,
16    suspicious record types) a stream of high-improbability names under ONE
17    domain is the real tunneling tell, not any single name.
18 3. robust thresholds (median + kMAD, resistant to CDN/cloud hash outliers)
19    instead of mean3 on a non-normal distribution.
20 4. a data-driven allowlist of YOUR frequent domains, so CloudFront/S3/telemetry
21    hashes don't false-positive.
22
23 Report-only by design: a privacy gateway must never silently drop DNS on a
24 statistical trip (a false positive = "my internet broke"). It surfaces findings;
25 promoting a domain to an actual block stays a human/AdGuard-blocklist decision.
26
27 Footprint: single streaming pass (never loads the log into RAM); the model is a
28 ~4040 bigram table + a few floats + a small allowlist (a few KB on disk); per-
29 domain state is a handful of ints plus a subdomain set capped at 64.
30
31 Exit codes (scan): 0 = ran, clean 1 = ran, anomalies found 2 = FAILED to run
32 (no/corrupt model, unreadable or empty querylog) a failure is never reported as
33 "clean". selftest exits 3 if the detection logic itself is broken.
34
35 Usage:
36 dns_anomaly_detector.py learn [--log PATH] [--out PATH] # train from querylog
37 dns_anomaly_detector.py scan [--log PATH] [--model PATH] [--recent N] [--quiet]
38 dns_anomaly_detector.py selftest # prove it detects (exit 3 if broken)
39 dns_anomaly_detector.py demo # human-readable showcase
40
41 """
42 import argparse
43 import json
44 import math
45 import os
46 import random
47 import re
48 import sys
49 import time
50 from collections import defaultdict
51
52 DEFAULT_LOG = os.path.join(os.path.dirname(os.path.abspath(__file__)),
53                             "..", "adguard", "work", "data", "querylog.json")
54 DEFAULT_MODEL = os.path.join(os.path.dirname(os.path.abspath(__file__)),
55                               "..", ".dns_baseline.json")
56 OUTPUT_LOG = os.path.join(os.path.dirname(os.path.abspath(__file__)), "..", "output.log")
57
58 def log_fail(msg):
59     """Persist a hard failure to output.log (besides stderr) so a broken detector
60     is on the record even when run headless (launchd/sweep/CI). Heisenbug-safe by
61     construction: it only ever APPENDS to a separate file it touches neither
62     stdout/stderr (which sweep captures) nor any detection state, and a logging
63     error is swallowed so it can never change what the detector does or returns."""
64     try:
65         with open(OUTPUT_LOG, "a", encoding="utf-8") as f:
66             f.write(f"[{time.strftime('%Y-%m-%d %H:%M:%S')}] [dns-anomaly] FAIL: {msg}\n")
67     except OSError:
68         pass
69
70 ALPHABET = "abcdefghijklmnopqrstuvwxyz0123456789-_"
71 AIDX = {c: i for i, c in enumerate(ALPHABET)}
72 OTHER = len(ALPHABET) # bucket for any char outside the alphabet
73 NSYM = len(ALPHABET) + 1
74 BEG = NSYM # start-of-label marker
75 NSTATES = NSYM + 1
76
77 # Minimal two-level public suffixes so eTLD+1 grouping is stable without shipping
78 # the full ~200 KB Public Suffix List. Exact PSL isn't needed we only need a
79 # consistent grouping key per registered domain.
80 TWO_LEVEL = {
81     "co.uk", "org.uk", "ac.uk", "gov.uk", "co.jp", "co.kr", "co.in", "co.za",
82     "co.nz", "com.au", "net.au", "org.au", "com.br", "com.mx", "com.tr",
83     "com.cn", "com.hk", "com.sg", "com.tw", "com.ar", "co.il", "ne.jp",
84 }
85
86 def sym(ch):
87     return AIDX.get(ch, OTHER)
88

```

```

89
90
91 def registered_and_sub(qh):
92     """Split a query host into (registered_domain, subdomain_labels_string)."""
93     qh = qh.rstrip(".").lower()
94     parts = qh.split(".")
95     if len(parts) <= 2:
96         return qh, ""
97     last2 = ".".join(parts[-2:])
98     reg_n = 3 if last2 in TWO_LEVEL else 2
99     reg = ".".join(parts[-reg_n:])
100    sub = ".".join(parts[:-reg_n])
101    return reg, sub
102
103
104 def iter_querylog(path, recent=0):
105     """Stream (QH, QT) from an AdGuard querylog. recent>0 keeps only the last N."""
106     if recent > 0:
107         # Bounded tail without loading the file: keep a ring buffer of raw lines.
108         from collections import deque
109         buf = deque(maxlen=recent)
110         with open(path, "r", encoding="utf-8", errors="ignore") as f:
111             for line in f:
112                 if line.strip():
113                     buf.append(line)
114             lines = buf
115     else:
116         lines = _line_reader(path)
117     for line in lines:
118         try:
119             rec = json.loads(line)
120         except (json.JSONDecodeError, ValueError):
121             continue
122         qh = rec.get("QH")
123         if not qh:
124             continue
125         qh = qh.strip().lower()
126         if qh.endswith("in-addr.arpa") or qh.endswith("ip6.arpa"):
127             continue
128         yield qh, rec.get("QT", "")
129
130
131 def _line_reader(path):
132     with open(path, "r", encoding="utf-8", errors="ignore") as f:
133         for line in f:
134             if line.strip():
135                 yield line
136
137
138 # n-gram model
139 def label_improbability(label, logp):
140     """Mean -log2 P(char|prev) over a label. High = doesn't look like real naming."""
141     if not label:
142         return 0.0
143     prev = BEG
144     total = 0.0
145     for ch in label:
146         s = sym(ch)
147         total += logp[prev][s]
148         prev = s
149     return total / len(label)
150
151
152 def name_score(sub, logp):
153     """Worst-label improbability across the subdomain (tunneling hides in one label)."""
154     if not sub:
155         return 0.0
156     return max(label_improbability(lbl, logp) for lbl in sub.split("."))
157
158
159 def finalize_logp(counts):
160     """Add-1 smoothed transition matrix -log2 probabilities."""
161     logp = [[0.0] * NSYM for _ in range(NSTATES)]
162     for a in range(NSTATES):
163         row_total = sum(counts[a]) + NSYM # +NSYM for add-1 smoothing
164         for b in range(NSYM):
165             p = (counts[a][b] + 1) / row_total
166             logp[a][b] = -math.log2(p)
167     return logp
168
169

```

```

170 def learn(log_path, out_path):
171     counts = [[0] * NSYM for _ in range(NSTATES)]
172     reg_counts = defaultdict(int)
173     reservoir = [] # sampled subdomain strings for threshold calibration
174     RES_MAX = 6000
175     seen = 0
176
177     for qh, _qt in iter_querylog(log_path):
178         reg, sub = registered_and_sub(qh)
179         reg_counts[reg] += 1
180         if not sub:
181             continue
182         for lbl in sub.split("."):
183             prev = BEG
184             for ch in lbl:
185                 s = sym(ch)
186                 counts[prev][s] += 1
187                 prev = s
188         # Reservoir-sample subdomains (unbiased) for a robust score distribution.
189         seen += 1
190         if len(reservoir) < RES_MAX:
191             reservoir.append(sub)
192         else:
193             j = random.randint(0, seen - 1)
194             if j < RES_MAX:
195                 reservoir[j] = sub
196
197     if seen == 0:
198         print("learn: no subdomained queries found nothing to train on.", file=sys.stderr)
199         return 1
200
201     logp = finalize_logp(counts)
202
203     # Robust threshold from the score distribution: median + kMAD (MAD1.4826,
204     # so k=6 mean+4 but resistant to the CDN-hash outliers that wreck plain ).
205     scores = sorted(name_score(s, logp) for s in reservoir if s)
206     median = _median(scores)
207     mad = _median(sorted(abs(x - median) for x in scores)) or 1e-9
208     threshold = median + 6.0 * mad
209
210     # Data-driven allowlist: the user's frequent registered domains (their normal
211     # traffic, incl. their habitual CDNs) tunneling to a NEW domain isn't here.
212     top = sorted(reg_counts.items(), key=lambda kv: kv[1], reverse=True)
213     allow_cut = max(20, int(0.002 * seen)) # "frequent" = seen a lot for this net
214     allowlist = [r for r, c in top if c >= allow_cut][:400]
215
216     model = {
217         "version": 2,
218         "counts": counts, # compact ints; logp recomputed on load
219         "threshold": threshold,
220         "median": median,
221         "mad": mad,
222         "allowlist": allowlist,
223         "trained_on": seen,
224         "unique_registered": len(reg_counts),
225     }
226     with open(out_path, "w", encoding="utf-8") as f:
227         json.dump(model, f, separators=(",", ":"))
228     print(f"learn: trained on {seen} subdomained queries across "
229           f"{len(reg_counts)} registered domains.")
230     print(f"score median={median:.3f} MAD={mad:.3f} threshold={threshold:.3f}")
231     print(f"allowlist: {len(allowlist)} frequent domains {out_path} "
232           f"{os.path.getsize(out_path)} bytes")
233     return 0
234
235
236 def _median(xs):
237     n = len(xs)
238     if n == 0:
239         return 0.0
240     xs = sorted(xs)
241     m = n // 2
242     return xs[m] if n % 2 else (xs[m - 1] + xs[m]) / 2.0
243
244
245 # scan
246 SUSPICIOUS_QT = {"TXT", "NULL", "CNAME", "ANY"} # classic tunneling data channels
247 SUB_CAP = 64 # bound per-domain memory
248
249 # High-entropy-but-legit CDN/cloud suffixes: their hostnames LOOK random but are
250 # structured, not encoded payload. A tiny static safety net beside the data-driven

```

```

251 # frequency allowlist (which only catches domains queried a lot in THIS log).
252 CDN_SUFFIXES = (
253     "nflxvideo.net", "nflxso.net", "googleusercontent.com", "usercontent.goog",
254     "cloudfront.net", "akamai.net", "akamaihd.net", "akamaiedge.net", "edgekey.net",
255     "edgesuite.net", "fastly.net", "fbcdn.net", "yting.com", "ggpht.com", "gvt1.com",
256     "gvt2.com", "ie100.net", "azureedge.net", "cloudflare.net", "cloudflarestream.com",
257     "digitaloceanspaces.com", "b-cdn.net", "llnwd.net", "wac.edgecastcdn.net", "cdn77.org",
258 )
259
260 _HEX = re.compile(r"^[0-9a-f]{16,}$")
261 _B32 = re.compile(r"^[a-z2-7]{16,}$")
262
263
264 def looks_encoded(label):
265     """True if a label looks like an encoded payload (hex/base32/long base64),
266     NOT just a random-ish structured CDN name or a long dictionary word. This is
267     the strongest single discriminator between DNS tunneling and legitimate
268     high-entropy hostnames."""
269     if _HEX.match(label):
270         # 16+ hex chars
271         return True
272     # base32 alphabet [a-z2-7] fully overlaps lowercase letters, so a long word
273     # like "visualstudioexptteam" would match require digits, which real base32
274     # payload always carries (~19% of chars) but dictionary words don't.
275     if _B32.match(label) and sum(c in "234567" for c in label) >= 2:
276         return True
277     # long base64-ish: high char diversity AND several digits (excludes words).
278     if len(label) >= 20 and len(set(label)) >= 13 and sum(c.isdigit() for c in label) >= 3:
279         return True
280     return False
281
282 def is_allowlisted(reg, allow):
283     return reg in allow or any(reg == s or reg.endswith(".") + s for s in CDN_SUFFIXES)
284
285
286 def scan(log_path, model_path, recent, quiet):
287     # Fail LOUD, never silently: a corrupt/absent model or an unreadable log must
288     # NOT look like a clean result that would be a security detector reporting
289     # "all clear" while blind. Each failure prints to stderr and exits non-zero.
290     try:
291         with open(model_path, "r", encoding="utf-8") as f:
292             model = json.load(f)
293             counts = model["counts"]
294             threshold = float(model["threshold"])
295             if (not isinstance(counts, list) or len(counts) != NSTATES
296                 or any(not isinstance(r, list) or len(r) != NSYM for r in counts)):
297                 raise ValueError("bigram matrix has the wrong shape")
298     except (OSError, json.JSONDecodeError, KeyError, ValueError, TypeError) as e:
299         msg = (f"could not load a usable baseline model at {model_path} "
300             f"({type(e).__name__}: {e}) run 'learn' first")
301         print(f"scan: FAILED to {msg}.", file=sys.stderr)
302         log_fail(msg)
303         return 2
304     logp = finalize_logp(counts)
305     allow = set(model.get("allowlist", []))
306
307     agg = {} # reg -> [queries, {subdomains cap64}, sum_score, n_scored, susp_qt, maxlen, encoded]
308     for qh, qt in iter_querylog(log_path, recent=recent):
309         reg, sub = registered_and_sub(qh)
310         a = agg.get(reg)
311         if a is None:
312             a = agg[reg] = [0, set(), 0.0, 0, 0, 0, 0]
313         a[0] += 1
314         if qt in SUSPICIOUS_QT:
315             a[4] += 1
316         if sub:
317             if len(a[1]) < SUB_CAP:
318                 a[1].add(sub)
319             sc = name_score(sub, logp)
320             a[2] += sc
321             a[3] += 1
322             for l in sub.split("."):
323                 if len(l) > a[5]:
324                     a[5] = len(l)
325                 if not a[6] and looks_encoded(l):
326                     a[6] = 1
327
328     if not agg:
329         msg = (f"the querylog at {log_path} yielded 0 parseable DNS records "
330             f"the detector did NOT run; this is NOT a clean result")
331         print(f"scan: FAILED {msg}.", file=sys.stderr)

```

```

332     log_fail(msg)
333     return 2
334
335 findings = []
336 for reg, (q, subs, ssum, sn, susp, maxlen, encoded) in agg.items():
337     if sn == 0 or is_allowlisted(reg, allow):
338         continue
339     mean_score = ssum / sn
340     uniq = len(subs)
341     improbable = mean_score > threshold
342     # Three independent tunneling signatures (any one flags), each tunneling-
343     # SHAPED so structured CDN names don't trip it. Encoded-payload volume is
344     # primary because encoded labels score like CDN hashes under the bigram
345     # model the improbability threshold alone can't separate them, but their
346     # SHAPE (pure hex/base32/base64) and per-domain multiplicity can:
347     strong_encoded = encoded and uniq >= 8 # many unique encoded payloads
348     txt_abuse = susp >= 5 # sustained TXT/NULL channel
349     improbable_vol = improbable and (uniq >= 8 or maxlen >= 24)
350     if strong_encoded or txt_abuse or improbable_vol:
351         corro = min(1.0, uniq / 30.0 + susp / 12.0 + encoded * 0.45 + (maxlen >= 40) * 0.2)
352         over = min(1.0, max(0.0, mean_score - threshold) / (threshold + 1e-9))
353         findings.append((round(0.55 * corro + 0.45 * over, 3), reg, round(mean_score, 2),
354             uniq + (SUB_CAP if uniq >= SUB_CAP else 0), susp, maxlen, q))
355
356 findings.sort(reverse=True)
357 if quiet:
358     print(f"DNS anomaly scan: {len(findings)} suspicious domain(s) "
359         f"(threshold={threshold:.2f}, {len(agg)} domains scanned)")
360 else:
361     print(f"\n=== DNS anomaly scan {len(findings)} suspicious domain(s) "
362         f"of {len(agg)} scanned (threshold {threshold:.2f}) ===")
363     if findings:
364         print(f"{'conf':>5} {'registered domain':32} {'score':>6} {'uniqusub':>7} "
365             f"{'txt/null':>8} {'maxlbl':>6} {'queries':>7}")
366         for conf, reg, sc, uniq, susp, maxlen, q in findings[:25]:
367             cap = "+" if uniq > SUB_CAP else " "
368             print(f"{'conf':>5} {'reg':32.32} {'sc':>6} {'min(uniq, SUB_CAP):>6}{cap} "
369                 f"{'susp':>8} {'maxlen':>6} {'q':>7}")
370     else:
371         print(" clean no domain combined improbable naming with a volumetric signal.")
372 return 1 if findings else 0
373
374 # demo (self-test, no querylog needed)
375 def demo():
376     normal = ["google.com", "api.github.com", "mail.google.com", "cdn.jsdelivr.net",
377         "en.wikipedia.org", "outlook.office365.com", "i.reddit.it", "apple.com",
378         "static.cloudflareinsights.com", "play.googleapis.com"] * 40
379     counts = [[0] * NSYM for _ in range(NSTATES)]
380     for qh in normal:
381         _reg, sub = registered_and_sub(qh)
382         for lbl in sub.split("."):
383             prev = BEG
384             for ch in lbl:
385                 s = sym(ch)
386                 counts[prev][s] += 1
387                 prev = s
388     logp = finalize_logp(counts)
389     samples = [registered_and_sub(q)[1] for q in normal]
390     scores = sorted(name_score(s, logp) for s in samples if s)
391     med = _median(scores)
392     mad = _median(sorted(abs(x - med) for x in scores)) or 1e-9
393     thr = med + 6.0 * mad
394     print(f"demo threshold = {thr:.3f} (median {med:.3f}, MAD {mad:.3f})")
395     tests = [("api.github.com", "normal"), ("mail.google.com", "normal"),
396         ("A8F93BD72Q9X.attacker.com", "base64 tunnel"),
397         ("c732488a09b2e4f6d1.tunnel.evil.org", "hex tunnel"),
398         ("kZ9x-Qw82-Lp01-Vn55.dns.io", "high-entropy tunnel")]
399     for qh, label in tests:
400         _reg, sub = registered_and_sub(qh)
401         sc = name_score(sub, logp)
402         flag = "anomalous" if sc > thr else "normal"
403         print(f" {flag:14} score={sc:5.2f} {qh} ({label})")
404
405
406 def selftest():
407     """Prove the detection logic works end-to-end. Exit 0 = healthy, 3 = BROKEN
408     (loud). Lets sweep/CI catch a regression instead of silently trusting a
409     detector that no longer detects. Validates BOTH signal paths: the n-gram
410     language model (improbability) and the encoding-shape detector."""
411     normal = ["google.com", "api.github.com", "mail.google.com", "cdn.jsdelivr.net",

```



```

413         "en.wikipedia.org", "outlook.office365.com", "i.redd.it", "apple.com"] * 40
414 counts = [[0] * NSYM for _ in range(NSTATES)]
415 for qh in normal:
416     for lbl in registered_and_sub(qh)[1].split("."):
417         prev = BEG
418         for ch in lbl:
419             s = sym(ch)
420             counts[prev][s] += 1
421             prev = s
422 logp = finalize_logp(counts)
423 scores = sorted(name_score(registered_and_sub(q)[1], logp) for q in normal
424                 if registered_and_sub(q)[1])
425 med = _median(scores)
426 mad = _median(sorted(abs(x - med) for x in scores)) or 1e-9
427 thr = med + 6.0 * mad
428
429 fails = []
430 for g in ("mail.google.com", "api.github.com", "cdn.jsdelivr.net"):
431     sub = registered_and_sub(g)[1]
432     if name_score(sub, logp) > thr or any(looks_encoded(l) for l in sub.split(".")):
433         fails.append(f"false-positive on known-good '{g}'")
434 # hex/base32 encoded payloads must be caught by the SHAPE detector:
435 for enc in ("a8f93bd72c9e4f16.evil.com", "mzxw6ytb2do4zb3q7a.tunnel.io"):
436     lbl = registered_and_sub(enc)[1].split(".")[0]
437     if not looks_encoded(lbl):
438         fails.append(f"encoding detector missed '{lbl}'")
439 # an unpronounceable non-encoded label must be caught by the n-gram model:
440 imp = registered_and_sub("xqzjvkwbmqpzrfrxn.dga.net")[1].split(".")[0]
441 if name_score(imp, logp) <= thr:
442     fails.append(f"n-gram model missed improbable label '{imp}'")
443
444 if fails:
445     print("SELFTEST FAILED detector is not working:", file=sys.stderr)
446     for f in fails:
447         print(f" - {f}", file=sys.stderr)
448     log_fail("selftest detector is not working: " + "; ".join(fails))
449     return 3
450 print("selftest OK flags hex/base32 payloads + improbable labels, clears known-good.")
451 return 0
452
453
454 def main():
455     ap = argparse.ArgumentParser(description="nullexit DNS anomaly detector (report-only)")
456     sub = ap.add_subparsers(dest="cmd", required=True)
457     pl = sub.add_parser("learn"); pl.add_argument("--log", default=DEFAULT_LOG); pl.add_argument("--out",
458         default=DEFAULT_MODEL)
459     ps = sub.add_parser("scan")
460     ps.add_argument("--log", default=DEFAULT_LOG); ps.add_argument("--model", default=DEFAULT_MODEL)
461     ps.add_argument("--recent", type=int, default=0, help="scan only the last N records (0=all)")
462     ps.add_argument("--quiet", action="store_true")
463     sub.add_parser("demo")
464     sub.add_parser("selftest")
465     args = ap.parse_args()
466
467     if args.cmd == "learn":
468         if not os.path.exists(args.log):
469             print(f"learn: querylog not found at {args.log}", file=sys.stderr); return 1
470         return learn(args.log, args.out)
471     if args.cmd == "scan":
472         if not os.path.exists(args.log):
473             print(f"scan: querylog not found at {args.log}", file=sys.stderr); return 2
474         return scan(args.log, args.model, args.recent, args.quiet)
475     if args.cmd == "demo":
476         demo(); return 0
477     if args.cmd == "selftest":
478         return selftest()
479
480 if __name__ == "__main__":
481     sys.exit(main())

```

39 scripts/fix-docker-bridge-collision.sh

```

1 #!/bin/bash
2 # scripts/fix-docker-bridge-collision.sh nullexit fix for devref 9.13
3 #

```

```

4 # Symptom: host default route points at Docker's bridge gateway (172.17.0.1)
5 # instead of the real Wi-Fi gateway. ALL host internet egress fails. Remote
6 # tailnet clients route through the gateway correctly (their tunnel is
7 # independent of the host's broken route table), so they look healthy while
8 # the host itself cannot reach anything.
9 #
10 # Root cause: Docker's default bridge claims 172.17.0.0/16 by default. When the
11 # host's physical Wi-Fi also lands in the same /16 extremely common on
12 # university networks which hand out 172.16.0.0/12 the kernel
13 # installs Docker's bridge as the default route, and every packet bound for
14 # the internet is silently absorbed by docker0 instead of egressing.
15 #
16 # This script:
17 # 1. Detects the actual collision (Docker's 172.17.0.0/16 vs the host's
18 # Wi-Fi subnet, queried live via the ACTIVE interface not hardcoded en0)
19 # 2. Picks a non-colliding Docker bip using /8-family routing (works for
20 # 172.16.0.0/12 campus networks; naive prefix-match would pick 172.26/24
21 # which still lives INSIDE a /12 of 172.16.0.0/12)
22 # 3. Backs up ~/.colima/default/colima.yaml with an ISO-8601 timestamp
23 # 4. Edits the file in place, idempotent (replace existing docker.bip /
24 # append new docker.bip without disturbing other keys; skip YAML comments)
25 # 5. Restarts Colima so the VM rebuilds with the new bridge subnet
26 # 6. Rebinds Wi-Fi DHCP so the host re-learns its real Wi-Fi gateway
27 # 7. Verifies the new default route is no longer the Docker bridge
28 # 8. Re-runs toggle.sh to bring the gateway back up cleanly
29 # 9. Re-runs scripts/diagnose-host-leak.sh and prints the new verdict
30 #
31 # Usage:
32 # bash scripts/fix-docker-bridge-collision.sh # apply the fix
33 # bash scripts/fix-docker-bridge-collision.sh --dry # show what would change, then exit
34 # bash scripts/fix-docker-bridge-collision.sh --skip-toggle # fix but don't restart gateway
35 # bash scripts/fix-docker-bridge-collision.sh --help # usage
36
37 set -uo pipefail
38 # NOTE: errexit ('set -e') is intentionally NOT enabled so intermediate
39 # 'command || warn' patterns can continue past non-fatal failures.
40 # Every critical failure (colima restart, sed, cp backup, toggle.sh)
41 # uses an explicit '|| die' so a half-applied state never escapes.
42
43 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
44 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
45 LOG_FILE="$PROJECT_ROOT/output.log"
46 source "$SCRIPT_DIR/common.sh"
47 TIMESTAMP="$(date -u +%Y%m%dT%H%M%S)"
48 COLIMA_YAML="$HOME/.colima/default/colima.yaml"
49 PLATFORM="$(uname -s)"
50
51 # Argument parsing
52 DRY_RUN=false
53 SKIP_TOGGLE=false
54 case "${1:-}" in
55     --dry|-n) DRY_RUN=true ;;
56     --skip-toggle) SKIP_TOGGLE=true ;;
57     --help|-h)
58         sed -n '28,37p' "$0"
59         exit 0
60     ;;
61     *)
62         : # default: apply
63     ;;
64 *)
65     printf 'Unknown argument: %s\n' "$1" >&2
66     exit 2
67 ;;
68 esac
69
70
71
72 # Helper: print "[dry-run] $cmd" or actually run $cmd. Wraps the redirect
73 # INSIDE the function so outer '>> $FILE' statements never accidentally
74 # append dry-run noise to the real file. (Earlier draft didn't and the
75 # dry-run polluted $COLIMA_YAML.)
76 run() {
77     if [ "$DRY_RUN" = "true" ]; then
78         printf '[dry-run] %s\n' "$*"
79         return 0
80     fi
81     printf '%s\n' "$*"
82     "$@"
83 }
84

```

```

85 # Platform-aware sed in-place. macOS BSD sed needs an empty argument;
86 # GNU sed (Linux) does not. Earlier draft used '-i '' everywhere and
87 # would have failed on Linux silently.
88 sed_inplace() {
89     local expr="$1"
90     local file="$2"
91     case "$PLATFORM" in
92         Darwin) run sed -i '' "$expr" "$file" ;;
93         *)      run sed -i "$expr" "$file" ;;
94     esac
95 }
96
97 # 0. Resolve the ACTIVE physical interface once (used everywhere)
98 # Replaces the earlier 'en0' hardcoding. Same algorithm recover.sh uses:
99 # 'route get default' if utun/tun, walk en0..en5 first with 'status: active'
100 # + 'inet' fall back to en0. This means the script also works on hosts
101 # whose primary NIC is en1/en2 (Thunderbolt Ethernet, USB-LAN, etc.).
102 resolve_active_iface() {
103     local iface
104     iface=$(route get default 2>> "$LOG_FILE" | awk '/interface:/{print $2; exit}')
105     if [ -z "$iface" ] || [ "${iface#tun}" != "$iface" ] || [ "${iface#tun}" != "$iface" ]; then
106         for i in en0 en1 en2 en3 en4 en5; do
107             if ifconfig "$i" 2>> "$LOG_FILE" | grep -q "status: active" \
108                 && ifconfig "$i" 2>> "$LOG_FILE" | grep -q "inet"; then
109                 iface="$i"; break
110             fi
111         done
112     fi
113     [ -z "$iface" ] && iface="en0"
114     printf '%s\n' "$iface"
115 }
116 ACTIVE_IFACE=$(resolve_active_iface)
117
118 # 1. Detect collision
119 step "1/9 Detecting Docker-bridge vs Wi-Fi subnet collision"
120
121 # Host subnet on the ACTIVE interface (not hardcoded en0).
122 HOST_SUBNET=$(ifconfig "$ACTIVE_IFACE" 2>> "$LOG_FILE" \
123     | awk '/inet /{print $2; exit}')
124 HOST_GATEWAY=$(route get default 2>> "$LOG_FILE" \
125     | awk '/gateway:/{print $2; exit}')
126 DEFAULT_DEV=$(route get default 2>> "$LOG_FILE" \
127     | awk '/interface:/{print $2; exit}')
128
129 printf ' active iface: %s\n' "$ACTIVE_IFACE"
130 printf '  iface inet: %s\n' "${HOST_SUBNET:-none}"
131 printf ' default via: %s\n' "${HOST_GATEWAY:-none}"
132 printf ' default dev: %s\n' "${DEFAULT_DEV:-none}"
133
134 # Collision fires when the host's default-route gateway is 172.17.0.1 (Docker's
135 # bridge) OR when the host subnet is in 172.17.0.0/16. The full-bridge-collision
136 # case (gateway is literally the docker bridge IP) is the smoking gun.
137 COLLISION=false
138 if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#172.17.}" != "$HOST_GATEWAY" ]; then
139     COLLISION=true
140     fail "default gateway $HOST_GATEWAY collides with Docker's default bridge (172.17.0.0/16)"
141 elif [ -n "$HOST_SUBNET" ] && [ "${HOST_SUBNET#172.17.}" != "$HOST_SUBNET" ]; then
142     COLLISION=true
143     fail "host subnet $HOST_SUBNET overlaps Docker's 172.17.0.0/16 bridge"
144 else
145     ok "no collision detected between host Wi-Fi and Docker's default bridge"
146 fi
147
148 if [ "$COLLISION" = "false" ]; then
149     echo ""
150     echo "Nothing to fix. If you still see egress issues, run:"
151     echo "    bash scripts/diagnose-host-leak.sh"
152     exit 0
153 fi
154
155 # 2. Pick non-colliding bip via /8-family routing
156 step "2/9 Picking a non-colliding Docker bip for ~/.colima/default/colima.yaml"
157
158 # Earlier draft tried naive prefix-match: "if candidate's 3-octet prefix is
159 # not a literal prefix of host gateway, pick it". That breaks on /12 campus
160 # networks (172.16.0.0/12 covers 172.16-172.31): a candidate like
161 # 172.26.0.1/24 LIVES INSIDE the host's /12 even though the prefix test
162 # passes. Fix: classify the host's address family FIRST, then choose ALL
163 # candidates from a DIFFERENT RFC1918 family. /8 boundaries are coarse
164 # enough that no /24 picked this way can ever collide.
165 case "$HOST_GATEWAY" in

```

```

166 172.16.*|172.17.*|172.18.*|172.19.*|172.20.*|172.21.*|172.22.*|172.23.*|\
167 172.24.*|172.25.*|172.26.*|172.27.*|172.28.*|172.29.*|172.30.*|172.31.*)
168 HOST_FAMILY="172"
169 # Host is in 172.x jump past 172.31 entirely. Pick from 10.x / 192.168.x.
170 CANDIDATES="10.200.0.1/24 10.201.0.1/24 192.168.99.1/24"
171 ;;
172 10.*)
173 HOST_FAMILY="10"
174 # Host is in 10.x pick from 172.x (well outside any commonly-deployed
175 # 10.0.0.0/8 slice) / 192.168.x.
176 CANDIDATES="172.26.0.1/24 172.28.0.1/24 192.168.99.1/24"
177 ;;
178 192.168.*)
179 HOST_FAMILY="192"
180 CANDIDATES="172.26.0.1/24 10.200.0.1/24 10.201.0.1/24"
181 ;;
182 *)
183 HOST_FAMILY="unknown"
184 warn "host gateway $HOST_GATEWAY is not in any RFC1918 range (/8 family)"
185 warn "falling back to abstract candidate set; verify collision-free manually"
186 CANDIDATES="172.26.0.1/24 10.200.0.1/24 192.168.99.1/24"
187 ;;
188 esac
189
190 printf ' host address family: %s\n' "$HOST_FAMILY"
191 printf ' candidate pool: %s\n' "$CANDIDATES"
192
193 # Sanity-check the chosen pool has at least one entry that doesn't literally
194 # share a /24 with the host gateway. Naive check (just /24 prefix) is fine
195 # here because we already picked across families.
196 CHOSEN=""
197 for cand in $CANDIDATES; do
198 PREFIX3=$(printf '%s' "$cand" | cut -d. -f1-3)
199 if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#"$PREFIX3"}" != "$HOST_GATEWAY" ]; then
200 warn "skip $cand (literal /24 collision with host gateway $HOST_GATEWAY)"
201 continue
202 fi
203 CHOSEN="$cand"
204 ok "chose bip=$CHOSEN (cross-family, no /24 literal collision)"
205 break
206 done
207
208 if [ -z "$CHOSEN" ]; then
209 CHOSEN="172.26.0.1/24"
210 warn "could not find a clean candidate; defaulting to $CHOSEN"
211 warn "verify with 'route get default' after restart"
212 fi
213
214 # 3. Backup colima.yaml
215 step "3/9 Backing up $COLIMA_YAML"
216 mkdir -p "$(dirname "$COLIMA_YAML")"
217 BACKUP="$HOME/.colima/default/colima.yaml.bak.$TIMESTAMP"
218 if [ -f "$COLIMA_YAML" ]; then
219 # 'cp' is safe in dry-run because run() suppresses the inner command
220 # entirely no chance of accidentally writing a real backup.
221 run cp -p "$COLIMA_YAML" "$BACKUP" || die "backup failed"
222 if [ "$DRY_RUN" = "false" ]; then ok "backup: $BACKUP"; fi
223 else
224 warn "no existing $COLIMA_YAML (will create fresh)"
225 fi
226
227 # 4. Edit colima.yaml in place (idempotent)
228 step "4/9 Setting docker.bip=$CHOSEN in colima.yaml (in place)"
229
230 # Resolve the existing docker.bip (if any) without picking up commented-out
231 # YAML keys. awk-based extraction; the negated-comment check (! /^[[:space:]]*#')
232 # ensures '# bip: 1.2.3.4/24' is treated as a comment, not the active value.
233 EXISTING_BIP=""
234 if [ -f "$COLIMA_YAML" ]; then
235 EXISTING_BIP=$(awk '
236 /^[[:space:]]*#/{next}
237 /^[[:space:]]*bip:[[:space:]]*/{print $2; exit}
238 ' "$COLIMA_YAML" 2>> "$LOG_FILE" || true)
239 fi
240
241 if [ -n "$EXISTING_BIP" ] && [ "$EXISTING_BIP" = "$CHOSEN" ]; then
242 ok "colima.yaml already has docker.bip=$CHOSEN no edit needed"
243 elif [ -n "$EXISTING_BIP" ]; then
244 printf ' (replacing existing docker.bip=%s with %s)\n' "$EXISTING_BIP" "$CHOSEN"
245 # Use sed_inplace so Linux / macOS get the right syntax. Skip YAML comment
246 # lines so existing '# bip:' comments aren't rewired.

```

```

247 sed_inplace "s|^\\([[:space:]]*bip:[[:space:]]*\\).*\\$|\\1$CHOSEN|" "$COLIMA_YAML" \
248 || die "sed replacement failed (colima.yaml malformed?)"
249 ok "colima.yaml updated"
250 else
251 # No existing 'bip:' key. Branch:
252 # - file has 'docker:' block append ' bip:' under it
253 # - file has no 'docker:' block append new docker: section
254 # - file does not exist create fresh
255 # Each branch is dry-run-aware so no branch can accidentally leak noise
256 # into $COLIMA_YAML during a preview.
257 HAS_DOCKER_BLOCK=false
258 if [ -f "$COLIMA_YAML" ] && grep -q '^docker:' "$COLIMA_YAML"; then
259 HAS_DOCKER_BLOCK=true
260 fi
261
262 if [ "$HAS_DOCKER_BLOCK" = "true" ]; then
263 if [ "$DRY_RUN" = "true" ]; then
264 printf ' [dry-run] would append to %s:\\n bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
265 else
266 printf '\\n bip: %s\\n' "$CHOSEN" >> "$COLIMA_YAML" \
267 || die "append failed"
268 ok "appended bip under existing docker: block"
269 fi
270 elif [ -f "$COLIMA_YAML" ]; then
271 if [ "$DRY_RUN" = "true" ]; then
272 printf ' [dry-run] would append to %s:\\ndocker:\\n bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
273 else
274 printf '\\ndocker:\\n bip: %s\\n' "$CHOSEN" >> "$COLIMA_YAML" \
275 || die "append failed"
276 ok "appended new docker: block at end of file"
277 fi
278 else
279 if [ "$DRY_RUN" = "true" ]; then
280 printf ' [dry-run] would create %s with:\\ndocker:\\n bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
281 else
282 printf 'docker:\\n bip: %s\\n' "$CHOSEN" > "$COLIMA_YAML" \
283 || die "create failed"
284 ok "created fresh $COLIMA_YAML with docker.bip=$CHOSEN"
285 fi
286 fi
287 fi
288
289 # Always show the final colima.yaml docker section for visibility
290 printf '\\n%b colima.yaml 'docker:' section now reads:%b\\n' "$BOLD" "$NC"
291 awk '
292 /^docker:/{p=1}
293 p && /^[^ ]/ && !/^docker:/{exit}
294 p
295 ' "$COLIMA_YAML" 2>/dev/null | sed 's/^/ /' || true
296
297 if [ "$DRY_RUN" = "true" ]; then
298 printf '\\n %b(dry-run: exiting before colima restart / dhcp rebind / toggle.sh)%b\\n' "$YELLOW" "$NC"
299 exit 0
300 fi
301
302 # 5. Restart Colima
303 step "5/9 Restarting Colima VM with new bip"
304 if ! command -v colima >> "$LOG_FILE" 2>&1; then
305 die "colima not on PATH install with 'brew install colima' first"
306 fi
307 # No timeout wrapper (would mask actual VM startup time). If colima hangs,
308 # the user kills with Ctrl+C and can manually 'colima restart' from a fresh
309 # terminal the rest of this script is idempotent so re-running works.
310 colima restart >> "$LOG_FILE" 2>&1 || die "colima restart failed; check $LOG_FILE"
311 ok "colima restarted"
312
313 # 6. Rebind DHCP on the ACTIVE interface (not hardcoded en0)
314 step "6/9 Rebinding DHCP on $ACTIVE_IFACE so host re-learns its real gateway"
315
316 # Map ifname macOS network service name (Wi-Fi, USB 10/100 LAN, etc.).
317 # Same logic as the diagnostic + toggle.sh.
318 ACTIVE_SVC=""
319 if [ "$PLATFORM" = "Darwin" ]; then
320 ACTIVE_SVC=$(get_active_service)
321 fi
322
323 case "$PLATFORM" in
324 Darwin)
325 sudo -n networksetup -setdhcp "$ACTIVE_SVC" renew >> "$LOG_FILE" 2>&1 \
326 && ok "DHCP rebind issued on $ACTIVE_SVC" \
327 || warn "setdhcp renew failed; try 'sudo ipconfig set $ACTIVE_IFACE DHCP' manually"

```

```

328 ;;
329 Linux)
330     sudo -n dhclient -r "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 || true
331     sudo -n dhclient "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 \
332         && ok "dhclient rebind on $ACTIVE_IFACE" \
333         || warn "dhclient rebind failed"
334 ;;
335 *)
336     warn "unsupported platform $PLATFORM; skip DHCP rebind"
337 ;;
338 esac
339
340 # 7. Verify new default route
341 step "7/9 Verifying the new default route is NOT the Docker bridge"
342 sleep 2 # give DHCP + kernel a beat
343 NEW_DEFAULT=$(netstat -rn 2>> "$LOG_FILE" | awk '/^default/ {print $0; exit}')
344 printf '%s\n' "${NEW_DEFAULT:-no default route detected}"
345 NEW_GATEWAY=$(printf '%s' "$NEW_DEFAULT" | awk '{print $2; exit}')
346
347 if [ -n "$NEW_GATEWAY" ] && [ "${NEW_GATEWAY#172.17.}" != "$NEW_GATEWAY" ]; then
348     fail "default route STILL points at 172.17.x.x DHCP rebind may need manual touch"
349     warn "try: sudo ipconfig set $ACTIVE_IFACE DHCP (or reconnect Wi-Fi in the menu bar)"
350 elif [ -n "$NEW_GATEWAY" ]; then
351     ok "default route is now via $NEW_GATEWAY (no longer Docker's bridge)"
352 else
353     warn "could not determine new gateway; verify with 'route get default'"
354 fi
355
356 # 8. Re-run toggle.sh to bring the gateway back up
357 if [ "$SKIP_TOGGLE" = "true" ]; then
358     step "8/9 Skipping toggle.sh (--skip-toggle)"
359     warn "gateway is currently down run ./toggle.sh manually when ready"
360 else
361     step "8/9 Re-running toggle.sh to bring the gateway back up"
362     if bash "$PROJECT_ROOT/toggle.sh" 2>&1 | tee -a "$LOG_FILE"; then
363         ok "toggle.sh completed"
364     else
365         warn "toggle.sh returned non-zero see $LOG_FILE for details"
366     fi
367 fi
368
369 # 9. Re-run the diagnostic + show verdict
370 step "9/9 Re-running host-leak diagnostic to confirm fix"
371 if bash "$SCRIPT_DIR/diagnose-host-leak.sh" 2>&1 | tee -a "$LOG_FILE"; then
372     LATEST_REPORT=$(ls -t "$PROJECT_ROOT"/host-leak-diagnostic-*.txt 2>/dev/null | head -1)
373     if [ -n "$LATEST_REPORT" ]; then
374         VERDICT=$(awk '/^ Scenario:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
375         WARP=$(awk -F': *' '/^ WARP egress:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
376         printf '\n Most recent verdict: Scenario=%s WARP=%s\n' \
377             "${VERDICT:-?}" "${WARP:-?}"
378         if [ "$VERDICT" = "OK" ]; then
379             ok "fix confirmed by diagnostic host traffic is going through WARP"
380         fi
381     fi
382 else
383     warn "diagnostic exited non-zero; verify manually"
384 fi
385
386 printf '\n%b%b\n' "$BOLD" "$NC"
387 printf '%bDone.%b Full transcript at %s\n' "$BOLD" "$NC" "$LOG_FILE"
388 printf '%b%b\n' "$BOLD" "$NC"
389 exit 0

```

40 scripts/fix-ssh-delay.sh

```

1 #!/bin/bash
2 # scripts/fix-ssh-delay.sh Fix ~20s SSH connection delay caused by UseDNS
3 #
4 # macOS OpenSSH defaults UseDNS to "yes". When Tailscale SSH connects from
5 # a peer (phone/laptop), sshd performs a reverse-DNS (PTR) lookup on the
6 # client's 100.x.x.x Tailscale IP. Since DNS is hijacked to the nullexit
7 # gateway AdGuard Cloudflare WARP, and 100.64.0.0/10 is CGNAT private
8 # space, the PTR query times out (~15-20s) before SSH proceeds.
9 #
10 # This script uncomments "UseDNS no" in /etc/ssh/sshd_config and reloads
11 # sshd. Existing SSH sessions are NOT dropped.

```

```

12 #
13 # Usage:
14 #   sudo bash scripts/fix-ssh-delay.sh
15
16 set -euo pipefail
17
18 SSHD_CONFIG="/etc/ssh/sshd_config"
19
20 # Check we're root
21 if [ "$(id -u)" -ne 0 ]; then
22     echo "ERROR: This script must be run with sudo (it edits $SSHD_CONFIG)"
23     echo "Usage: sudo bash scripts/fix-ssh-delay.sh"
24     exit 1
25 fi
26
27 # Apply config (idempotent)
28 if grep -qE '^UseDNS no' "$SSHD_CONFIG" 2>/dev/null; then
29     echo "UseDNS is already set to 'no' in $SSHD_CONFIG"
30 elif grep -qE '^#UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
31     echo "Uncommenting 'UseDNS no' in $SSHD_CONFIG"
32     sed -i '' 's/^#UseDNS no/UseDNS no/' "$SSHD_CONFIG"
33 elif grep -qEi '^UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
34     echo "Changing existing UseDNS line to 'UseDNS no'"
35     sed -i '' 's/^UseDNS.*/UseDNS no/' "$SSHD_CONFIG"
36 else
37     echo "Appending 'UseDNS no' to $SSHD_CONFIG"
38     echo "" >> "$SSHD_CONFIG"
39     echo "UseDNS no" >> "$SSHD_CONFIG"
40 fi
41
42 # Verify
43 if ! grep -qE '^UseDNS no' "$SSHD_CONFIG"; then
44     echo "Failed to apply UseDNS no check $SSHD_CONFIG manually"
45     exit 1
46 fi
47
48 # Reload sshd (does NOT drop existing connections)
49 echo "Reloading sshd..."
50
51 SSHD_PID=$(pgrep -f /usr/sbin/sshd 2>/dev/null | head -1 || true)
52 if [ -n "$SSHD_PID" ]; then
53     kill -HUP "$SSHD_PID"
54     echo "Sent SIGHUP to sshd (PID $SSHD_PID) config reloaded"
55 else
56     echo "No running sshd found. macOS starts sshd on-demand per connection."
57     echo "The new 'UseDNS no' config will take effect on the next SSH connection."
58 fi
59
60 echo ""
61 echo "Done. SSH connections should now connect instantly."
62 echo "Existing sessions were not interrupted."

```

41 scripts/generate_tex.py

```

1 #!/usr/bin/env python3
2 import os
3
4 PROJECT_ROOT = os.path.abspath(os.path.join(os.path.dirname(__file__), '..'))
5
6 IGNORE_DIRS = {'.git', '.github', 'adguard', '__pycache__', 'Recover Gateway.app', 'Toggle Gateway.app',
7               'Sweep Gateway.app', 'nullexit-knowledge-graph'}
8 IGNORE_FILES = {'.env', 'nullexit_unified.tex', 'nullexit_unified.pdf', 'output.log', 'blocked.log', '.DS_Store',
9               '.lan_p2p_detected', '.signatures', '.gateway_ip', 'TUNNEL_FAILED_CLOSED.marker', '.host_ips',
10              'refactor.md', 'sweep.md', '.build_hash', '.rules_hash', 'knowledge-graph.html', '.dns_baseline.json'}
11 IGNORE_EXTS = {'.pdf', '.tex', '.png', '.jpg', '.jpeg', '.zip', '.tar', '.gz', '.log', '.aux', '.fls', '.fdb_latexmk',
12              '.out', '.toc', '.xdv', '.synctex(busy)'}
13 # Transient runtime reports (timestamped names) carry live egress/peer IPs never
14 # embed them in the published quine. Matched by filename prefix.
15 IGNORE_PREFIXES = ('sweep-', 'host-leak-diagnostic-')
16
17 # Exclude from code block inclusion, but keep in tree
18 SKIP_CONTENT_FILES = {'LICENSE'}
19
20 PRIORITY_FILES = [
21     'README.md',

```



```

18     'agent.md',
19     'devref.md',
20     'diagrams.md',
21     'toggle.sh',
22     'recover.sh',
23     'setup.sh',
24     'scripts/setup-linux.sh',
25     'docker-compose.yml',
26     'scripts/common.sh',
27     'scripts/watcher.sh',
28     'scripts/crypto.sh'
29 ]
30
31 def get_language(filename):
32     ext = os.path.splitext(filename)[1].lower()
33     if ext == '.sh' or ext == '.applescript':
34         return 'bash'
35     elif ext == '.py':
36         return 'Python'
37     elif ext == '.go':
38         return 'Go'
39     elif ext in {'.yml', '.yaml'}:
40         return 'bash'
41     elif ext == '.plist':
42         return 'XML'
43     elif ext == '.md':
44         return ''
45     return ''
46
47 def escape_tex(text):
48     return text.replace('_', '\\_')
49
50 def generate_tree(dir_path, prefix=""):
51     tree_str = ""
52     entries = sorted(os.listdir(dir_path))
53     entries = [e for e in entries if e not in IGNORE_DIRS and e not in IGNORE_FILES and os.path.splitext(
54         e)[1].lower() not in IGNORE_EXTS]
55
56     for i, entry in enumerate(entries):
57         path = os.path.join(dir_path, entry)
58         is_last = (i == len(entries) - 1)
59         connector = "-- " if is_last else "|-- "
60         tree_str += f"{prefix}{connector}{entry}\n"
61
62         if os.path.isdir(path):
63             extension = " " if is_last else "|"
64             tree_str += generate_tree(path, prefix + extension)
65
66     return tree_str
67
68 def main():
69     tex_path = os.path.join(PROJECT_ROOT, 'nullexit_unified.tex')
70
71     preamble = r"""
72 \documentclass{article}
73 \usepackage[utf8]{inputenc}
74 \usepackage[margin=1in]{geometry}
75 \usepackage{listings}
76 \usepackage{xcolor}
77 \usepackage{hyperref}
78 \usepackage{tocloft}
79
80 \definecolor{codegreen}{rgb}{0,0.6,0}
81 \definecolor{codegray}{rgb}{0.5,0.5,0.5}
82 \definecolor{codepurple}{rgb}{0.58,0,0.82}
83 \definecolor{backcolour}{rgb}{0.97,0.97,0.96}
84
85 \lstdefinestyle{mystyle}{
86     backgroundcolor=\color{backcolour},
87     commentstyle=\color{codegreen},
88     keywordstyle=\color{blue}\bfseries,
89     numberstyle=\tiny\color{codegray},
90     stringstyle=\color{codepurple},
91     basicstyle=\ttfamily\scriptsize,
92     breakatwhitespace=false,
93     breaklines=true,
94     captionpos=b,
95     keepspaces=true,
96     numbers=left,
97     numbersep=5pt,
98     showspaces=false,
99     showstringspaces=false,

```



```

176
177 if __name__ == '__main__':
178     main()

```

42 scripts/logger/go.mod

```

1 module nullexit/logger
2
3 go 1.22

```

43 scripts/logger/main.go

```

1 package main
2
3 import (
4     "bufio"
5     "encoding/json"
6     "fmt"
7     "io"
8     "log"
9     "os"
10    "regexp"
11    "strings"
12    "time"
13 )
14
15 const (
16     queryLogPath = "/adguard_work/data/querylog.json"
17     procKmsgPath = "/proc/kmsg"
18     outputLogPath = "/app/blocked.log"
19 )
20
21 func logMessage(msg string) {
22     timestamp := time.Now().Format("2006-01-02 15:04:05")
23     formattedMsg := fmt.Sprintf("%s %s\n", timestamp, msg)
24     fmt.Print(formattedMsg)
25
26     f, err := os.OpenFile(outputLogPath, os.O_APPEND|os.O_CREATE|os.O_WRONLY, 0644)
27     if err != nil {
28         log.Printf("Error writing to log file: %v\n", err)
29         return
30     }
31     defer f.Close()
32     f.WriteString(formattedMsg)
33 }
34
35 func tailFile(filepath string, out chan<- string) {
36     for {
37         if _, err := os.Stat(filepath); os.IsNotExist(err) {
38             time.Sleep(1 * time.Second)
39             continue
40         }
41         break
42     }
43
44     f, err := os.Open(filepath)
45     if err != nil {
46         log.Printf("Error opening file %s: %v\n", filepath, err)
47         return
48     }
49     defer func() {
50         if f != nil {
51             f.Close()
52         }
53     }()
54
55     // Seek to end
56     f.Seek(0, io.SeekEnd)
57     reader := bufio.NewReader(f)
58
59     for {
60         line, err := reader.ReadString('\n')

```

```

61 if err != nil {
62     if err == io.EOF {
63         stat, err := os.Stat(filepath)
64         if err == nil {
65             currentOffset, _ := f.Seek(0, io.SeekCurrent)
66             if stat.Size() < currentOffset {
67                 // File truncated or rotated
68                 f.Close()
69                 for {
70                     if _, err := os.Stat(filepath); os.IsNotExist(err) {
71                         time.Sleep(1 * time.Second)
72                         continue
73                     }
74                     break
75                 }
76                 f, _ = os.Open(filepath)
77                 reader = bufio.NewReader(f)
78                 continue
79             }
80         }
81         time.Sleep(500 * time.Millisecond)
82         continue
83     }
84     log.Printf("Error reading file %s: %v\n", filepath, err)
85     time.Sleep(1 * time.Second)
86     continue
87 }
88 out <- line
89 }
90 }
91
92 type adguardLogEntry struct {
93     IP      string `json:"IP"`
94     QH      string `json:"QH"`
95     QT      string `json:"QT"`
96     Result struct {
97         IsFiltered bool `json:"IsFiltered"`
98         Reason      int  `json:"Reason"`
99         Rules       []struct {
100             Text string `json:"Text"`
101         } `json:"Rules"`
102     } `json:"Result"`
103 }
104
105 func monitorDNS() {
106     logMessage("[System] Starting DNS block logger...")
107     lines := make(chan string)
108     go tailFile(queryLogPath, lines)
109
110     reasonMap := map[int]string{
111         2: "CustomRule",
112         3: "BlockList",
113         4: "SafeBrowsing",
114         5: "ParentalControl",
115         6: "SafeSearch",
116         7: "BlockedService",
117     }
118
119     for line := range lines {
120         var data adguardLogEntry
121         if err := json.Unmarshal([]byte(line), &data); err != nil {
122             continue
123         }
124
125         res := data.Result
126         if res.IsFiltered || (res.Reason != 0 && res.Reason != 1) {
127             qh := data.QH
128             if qh == "" {
129                 qh = "unknown"
130             }
131             qt := data.QT
132             if qt == "" {
133                 qt = "unknown"
134             }
135             clientIP := data.IP
136             if clientIP == "" {
137                 clientIP = "unknown"
138             }
139
140             ruleText := "unknown rule"
141             if len(res.Rules) > 0 && res.Rules[0].Text != "" {

```

```

142     ruleText = res.Rules[0].Text
143 }
144
145 reasonStr, ok := reasonMap[res.Reason]
146 if !ok {
147     reasonStr = fmt.Sprintf("Reason-%d", res.Reason)
148 }
149
150 logMessage(fmt.Sprintf("[DNS] Blocked %s (Type: %s) for client %s | Reason: %s | Rule: %s", qh, qt,
    clientIP, reasonStr, ruleText))
151 }
152 }
153 }
154
155 func monitorIPs() {
156     logMessage("[System] Starting IP block logger...")
157
158     prefixRe := regexp.MustCompile(`(IP_BLOCK_[A-Z_]+):`)
159     srcRe := regexp.MustCompile(`SRC=([0-9a-fA-F\.:]+)`)
160     dstRe := regexp.MustCompile(`DST=([0-9a-fA-F\.:]+)`)
161     protoRe := regexp.MustCompile(`PROTO=([A-Z0-9]+)`)
162     sptRe := regexp.MustCompile(`SPT=(\d+)`)
163     dptRe := regexp.MustCompile(`DPT=(\d+)`)
164     inRe := regexp.MustCompile(`IN=([a-zA-Z0-9\.\-]+)`)
165     outRe := regexp.MustCompile(`OUT=([a-zA-Z0-9\.\-]+)`)
166
167     f, err := os.Open(procKmsgPath)
168     if err != nil {
169         if os.IsPermission(err) {
170             logMessage("[System] ERROR: Insufficient permissions to read /proc/kmsg. Enable CAP_SYSLOG.")
171         } else {
172             logMessage(fmt.Sprintf("[System] ERROR in IP logger: %v", err))
173         }
174         return
175     }
176     defer f.Close()
177
178     reader := bufio.NewReader(f)
179     for {
180         line, err := reader.ReadString('\n')
181         if err != nil {
182             if err == io.EOF {
183                 time.Sleep(100 * time.Millisecond)
184                 continue
185             }
186             logMessage(fmt.Sprintf("[System] ERROR reading kmsg: %v", err))
187             time.Sleep(1 * time.Second)
188             continue
189         }
190
191         if strings.Contains(line, "IP_BLOCK") {
192             match := prefixRe.FindStringSubmatch(line)
193             if len(match) > 1 {
194                 prefix := match[1]
195
196                 src := "unknown"
197                 if m := srcRe.FindStringSubmatch(line); len(m) > 1 {
198                     src = m[1]
199                 }
200
201                 dst := "unknown"
202                 if m := dstRe.FindStringSubmatch(line); len(m) > 1 {
203                     dst = m[1]
204                 }
205
206                 proto := "unknown"
207                 if m := protoRe.FindStringSubmatch(line); len(m) > 1 {
208                     proto = m[1]
209                 }
210
211                 spt := ""
212                 if m := sptRe.FindStringSubmatch(line); len(m) > 1 {
213                     spt = m[1]
214                 }
215
216                 dpt := ""
217                 if m := dptRe.FindStringSubmatch(line); len(m) > 1 {
218                     dpt = m[1]
219                 }
220
221                 inIf := ""

```

```

222     if m := inRe.FindStringSubmatch(line); len(m) > 1 {
223         inIf = m[1]
224     }
225
226     outIf := ""
227     if m := outRe.FindStringSubmatch(line); len(m) > 1 {
228         outIf = m[1]
229     }
230
231     direction := "inbound"
232     if strings.HasSuffix(prefix, "DST") {
233         direction = "outbound"
234     }
235
236     listType := "MALICIOUS"
237     if !strings.Contains(prefix, "MALICIOUS") {
238         parts := strings.Split(prefix, "_")
239         if len(parts) >= 3 {
240             listType = "GEO_" + parts[2]
241         }
242     }
243
244     portStr := ""
245     if dpt != "" {
246         portStr = ":" + dpt
247     }
248
249     srcPortStr := ""
250     if spt != "" {
251         srcPortStr = ":" + spt
252     }
253
254     ifStr := ""
255     if inIf != "" && outIf != "" {
256         ifStr = fmt.Sprintf("IF: %s->%s", inIf, outIf)
257     } else if inIf != "" || outIf != "" {
258         ifStr = fmt.Sprintf("IF: %s%s", inIf, outIf)
259     }
260
261     logMessage(fmt.Sprintf("[IP] Blocked %s to %s%s (Proto: %s) from %s%s (%s) | List: %s", direction
262 , dst, portStr, proto, src, srcPortStr, ifStr, listType))
263 }
264 }
265 }
266
267 func main() {
268     logMessage("[System] Nullexit Block Logger starting...")
269
270     go monitorDNS()
271     monitorIPs()
272 }

```

44 scripts/noise.sh

```

1  #!/bin/bash
2  # scripts/noise.sh  nullexit verifiable-random noise + cover-traffic padding
3  #
4  # OPT-IN and ACTIVE. Unlike sweep.sh (read-only), this module PUTS TRAFFIC ON
5  # THE WIRE. It stays completely inert unless NOISE_ENABLED=true in .env, so the
6  # default posture is unchanged and the sweep's read-only guarantee is intact.
7  #
8  # Two jobs, one primitive (CSPRNG bytes from the OS, /dev/urandom via os.urandom):
9  #
10 # 1. verify prove a freshly generated buffer is statistically random using
11 # three standard tests (Shannon entropy, NIST monobit, chi-square over the
12 # 256 byte values). This is the "we can verify it's truly random
13 # mathematically" requirement it is harmless and always allowed.
14 #
15 # 2. pad a cover-traffic daemon. It emits verified-random UDP datagrams
16 # toward the exit at a configured rate + jitter so a passive on-path
17 # observer (residence / campus / hostile network) sees a continuously
18 # varying encrypted flow instead of your real traffic envelope. This is
19 # the "dummy padding" defence the Threat Model flags as the countermeasure
20 # to metadata / traffic-flow correlation. It ONLY runs with NOISE_ENABLED=true.
21 #

```

```

22 # HONEST LIMITS: this is constant/random one-way padding, not a full
23 # traffic-analysis-resistant shaper. It blurs volume + timing on the uplink; it
24 # does not mimic a specific protocol, pad bidirectionally, or defeat a global
25 # active adversary. It costs bandwidth and (on battery devices) power hence
26 # opt-in, modest default rate. The DNS/IP anomaly detector is deliberately NOT
27 # here (left for later).
28 #
29 # Usage:
30 #   bash scripts/noise.sh verify [BYTES]      # randomness self-test (default 1 MiB)
31 #   bash scripts/noise.sh start                # launch the padding daemon (background)
32 #   bash scripts/noise.sh pad                 # run the padding loop in the foreground
33 #   bash scripts/noise.sh status               # is the daemon running? target + rate
34 #   bash scripts/noise.sh stop                # stop the daemon
35 #   bash scripts/noise.sh --help              # this message
36 #
37 # Persistence: install launchd/com.nullexit.noise.plist (see README) so padding
38 # auto-starts at boot/login. It is governed by the SAME NOISE_ENABLED switch no
39 # extra knob: the launchd job is a clean no-op whenever NOISE_ENABLED=false.
40 #
41 # .env knobs (all optional; sane defaults):
42 #   NOISE_ENABLED          master on/off switch (default false)
43 #   NOISE_PAD_TARGET       where to aim padding (default: the gateway Tailscale IP)
44 #   NOISE_PAD_PORT         UDP port on the target (default 9999; no listener needed)
45 #   NOISE_PAD_RATE_KBPS    target padding bitrate (default 64)
46 #   NOISE_PAD_PACKET_BYTES datagram payload size (default 1024)
47 #   NOISE_PAD_JITTER_MS    +/- timing jitter per packet (default 40)
48 #   NOISE_PAD_MODE         'constant' = always emit the target rate (default);
49 #                           'topup' = measure real egress each second and inject
50 #                           only the deficit, so the OBSERVED total stays flat
51 #                           whether idle or busy (stronger anti-correlation, no
52 #                           wasted bandwidth when you're already sending)
53 #   NOISE_PAD_PCT          set the target as a % of LINK CAPACITY (overrides
54 #                           NOISE_PAD_RATE_KBPS). % of capacity, NOT of current
55 #                           throughput which would drop padding to zero when idle
56 #   NOISE_LINK_MBPS        link capacity for NOISE_PAD_PCT (default 100)
57 #   NOISE_PAD_IFACE        interface topup measures egress on (default en0)
58
59 set -uo pipefail
60
61 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
62 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
63 cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
64 source "$SCRIPT_DIR/common.sh"
65 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
66
67 SELF="$SCRIPT_DIR/noise.sh"
68 NOISE_PID="/tmp/nullexit-noise.pid"
69 NOISE_LOG="$PROJECT_ROOT/output.log"
70
71 # .env-driven config
72 NOISE_ENABLED=$(read_env_var "NOISE_ENABLED" | tr '[:upper:]' '[:lower:]')
73 [ -z "$NOISE_ENABLED" ] && NOISE_ENABLED="false"
74
75 PAD_TARGET=$(read_env_var "NOISE_PAD_TARGET")
76 [ -z "$PAD_TARGET" ] && PAD_TARGET="$(read_adguard_ip)" # gateway Tailscale IP
77 PAD_PORT=$(read_env_var "NOISE_PAD_PORT"); [ -z "$PAD_PORT" ] && PAD_PORT=9999
78 PAD_RATE=$(read_env_var "NOISE_PAD_RATE_KBPS"); [ -z "$PAD_RATE" ] && PAD_RATE=64
79 PAD_BYTES=$(read_env_var "NOISE_PAD_PACKET_BYTES"); [ -z "$PAD_BYTES" ] && PAD_BYTES=1024
80 PAD_JITTER=$(read_env_var "NOISE_PAD_JITTER_MS"); [ -z "$PAD_JITTER" ] && PAD_JITTER=40
81 # 'constant' = always emit the target rate. 'topup' = measure real egress each
82 # second and inject only the deficit, so the OBSERVED total rate stays flat
83 # whether you're idle or busy (the strong anti-correlation property).
84 PAD_MODE=$(read_env_var "NOISE_PAD_MODE" | tr '[:upper:]' '[:lower:]'); [ -z "$PAD_MODE" ] && PAD_MODE=
85 constant
86 PAD_IFACE=$(read_env_var "NOISE_PAD_IFACE"); [ -z "$PAD_IFACE" ] && PAD_IFACE=en0
87 # Target rate: NOISE_PAD_PCT (a % of LINK CAPACITY a stable reference, NOT of
88 # current throughput, which would collapse padding to zero when idle) overrides
89 # the absolute NOISE_PAD_RATE_KBPS when set.
90 PAD_PCT=$(read_env_var "NOISE_PAD_PCT")
91 PAD_LINK_MBPS=$(read_env_var "NOISE_LINK_MBPS"); [ -z "$PAD_LINK_MBPS" ] && PAD_LINK_MBPS=100
92 if [ -n "$PAD_PCT" ]; then
93     PAD_RATE=$(awk "BEGIN{printf \"%.0f\\n\", ($PAD_PCT/100.0)*$PAD_LINK_MBPS*1000}")
94 fi
95
96 log() { echo "[$(date '+%Y-%m-%d %H:%M:%S')] [noise] $*" >> "$NOISE_LOG" 2>/dev/null || true; }
97
98 require_enabled() {
99     if [ "$NOISE_ENABLED" != "true" ]; then
100         echo "noise: disabled (NOISE_ENABLED != true in .env) refusing to emit traffic." >&2
101         echo "Set NOISE_ENABLED=true to arm cover-traffic padding." >&2
102         exit 3
103     fi
104 }

```



```

102 fi
103 }
104
105 # Randomness self-test (stdlib python3; no third-party deps)
106 # Exits 0 (PASS) / 1 (FAIL). Prints a human summary unless $2 == "quiet".
107 verify_random() {
108     local nbytes="${1:-1048576}" mode="${2:-loud}"
109     python3 - "$nbytes" "$mode" <<'PY'
110 import os, sys, math
111 from collections import Counter
112
113 n = int(sys.argv[1])
114 mode = sys.argv[2] if len(sys.argv) > 2 else "loud"
115 buf = os.urandom(n)
116 bits = n * 8
117
118 # 1) Shannon entropy over byte values (bits/byte, ideal = 8.0)
119 c = Counter(buf)
120 H = -sum((v / n) * math.log2(v / n) for v in c.values())
121
122 # 2) NIST monobit: 1-bit count should be ~half. z = (ones - bits/2)/sqrt(bits/4)
123 ones = sum(bin(b).count("1") for b in buf)
124 z = (ones - bits / 2) / math.sqrt(bits / 4)
125
126 # 3) Chi-square over the 256 byte values (df = 255, mean 255, sd ~22.6)
127 exp = n / 256.0
128 chi2 = sum((c.get(i, 0) - exp) ** 2 / exp for i in range(256))
129
130 # PASS windows chosen wide enough that true CSPRNG output passes ~always,
131 # yet a broken/patterned source (constant, low-entropy, biased) fails hard.
132 H_OK = H >= 7.99 if n >= 1 << 16 else H >= 7.5
133 Z_OK = abs(z) < 5.0
134 CHI_OK = 185.0 < chi2 < 345.0
135 ok = H_OK and Z_OK and CHI_OK
136
137 if mode != "quiet":
138     print(f"sample : {n} bytes ({bits} bits) from os.urandom (/dev/urandom CSPRNG)")
139     print(f"shannon H : {H:.5f} bits/byte (ideal 8.0) [{ 'ok' if H_OK else 'FAIL' }]")
140     print(f"monobit z : {z:.3f} (|z|<5, bias {ones/bits*100:.3f}% ones) [{ 'ok' if Z_OK else 'FAIL' }]")
141     print(f"chi-square : {chi2:.1f} (df=255, expect ~255, window 185-345) [{ 'ok' if CHI_OK else 'FAIL' }]")
142     print(f"verdict : {'PASS statistically random' if ok else 'FAIL not random enough'}")
143 sys.exit(0 if ok else 1)
144 PY
145 }
146
147 # The cover-traffic loop (foreground; 'start' backgrounds this)
148 pad_loop() {
149     require_enabled
150     # Own the PID file for BOTH run modes (manual 'start' and launchd '_boot'), so
151     # 'status'/'stop' behave identically however padding was launched. Cleared on exit.
152     echo $$ > "$NOISE_PID"
153     trap 'rm -f "$NOISE_PID"' EXIT INT TERM
154     # Gate: only ever emit noise we have just proven is statistically random.
155     if ! verify_random 1048576 quiet; then
156         echo "noise: randomness self-test FAILED refusing to emit non-random padding." >&2
157         log "pad aborted: randomness self-test failed"
158         exit 1
159     fi
160     log "pad start ${PAD_TARGET}:${PAD_PORT} mode=${PAD_MODE} rate=${PAD_RATE}kbps pkt=${PAD_BYTES}B
161         jitter=${PAD_JITTER}ms iface=${PAD_IFACE}"
162     python3 - "$PAD_TARGET" "$PAD_PORT" "$PAD_RATE" "$PAD_BYTES" "$PAD_JITTER" "$PAD_MODE" "$PAD_IFACE" <<'
163 PY'
164 import os, sys, socket, time, random, subprocess
165
166 target = sys.argv[1]
167 port = int(sys.argv[2])
168 target_bps = float(sys.argv[3]) * 1000.0 / 8.0 # kbps bytes/sec (target TOTAL on the wire)
169 pkt = int(sys.argv[4])
170 jitter = float(sys.argv[5]) / 1000.0 # ms sec
171 mode = sys.argv[6] # 'constant' | 'topup'
172 iface = sys.argv[7]
173
174 s = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
175
176 def obytes(ifc):
177     """Cumulative output bytes for an interface what an on-path observer sees
178     leave the NIC. Returns None if the counter can't be read (fail toward MORE cover)."""
179     try:

```

```

179     out = subprocess.run(["netstat", "-ibn", "-I", ifc],
180                           capture_output=True, text=True, timeout=2).stdout
181     for line in out.splitlines():
182         f = line.split()
183         # the "<Link#N>" row carries cumulative counters; Obytes is column 9
184         if len(f) >= 11 and f[0] == ifc and f[2].startswith("<Link"):
185             return int(f[9])
186     except Exception:
187         pass
188     return None
189
190
191 def emit(budget_bytes):
192     """Send ~budget_bytes of fresh CSPRNG padding, spread across ~1s with jitter.
193     Returns bytes actually sent (so topup can subtract its own contribution)."""
194     n = max(0, int(budget_bytes // pkt))
195     interval = 1.0 / n if n > 0 else 1.0
196     sent = 0
197     for _ in range(n):
198         try:
199             s.sendto(os.urandom(pkt), (target, port)); sent += pkt
200         except OSError:
201             pass # dropped is fine the bytes still left the NIC
202         time.sleep(max(0.0, interval + random.uniform(-jitter, jitter)))
203     return sent
204
205
206 try:
207     prev = obytes(iface) if mode == "topup" else None
208     pad_last = 0
209     while True:
210         t0 = time.time()
211         if mode == "topup":
212             now = obytes(iface)
213             if prev is not None and now is not None:
214                 total = max(0, now - prev) # all egress on iface this window
215                 real = max(0, total - pad_last) # minus our own last-window padding
216                 budget = max(0.0, target_bps - real)
217             else:
218                 budget = target_bps # counters unavailable full target
219             prev = now
220         else: # constant
221             budget = target_bps
222             pad_last = emit(budget)
223             rem = 1.0 - (time.time() - t0) # keep a ~1s accounting window
224             if rem > 0:
225                 time.sleep(rem)
226     except KeyboardInterrupt:
227         pass
228 PY
229 }
230
231 daemon_running() {
232     [ -f "$NOISE_PID" ] || return 1
233     local pid; pid="$(cat "$NOISE_PID" 2>/dev/null)"
234     [ -n "$pid" ] && kill -0 "$pid" 2>/dev/null
235 }
236
237 cmd_start() {
238     require_enabled
239     if daemon_running; then
240         echo "noise: padding daemon already running (pid $(cat "$NOISE_PID"))."
241         exit 0
242     fi
243     echo "noise: verifying randomness before arming padding"
244     verify_random 1048576 loud || { echo "noise: aborting self-test failed." >&2; exit 1; }
245     nohup "$SELF" pad >>"$NOISE_LOG" 2>&1 &
246     local npid=$!
247     sleep 0.3 # let pad_loop write its PID file so 'status' is immediately accurate
248     echo "noise: padding daemon armed (pid $npid) ${PAD_TARGET}:${PAD_PORT}, ${PAD_RATE} kbps."
249 }
250
251 cmd_stop() {
252     if daemon_running; then
253         local pid; pid="$(cat "$NOISE_PID)"
254         kill "$pid" 2>/dev/null && echo "noise: stopped padding daemon (pid $pid)."
255         log "pad stop (pid $pid)"
256     else
257         echo "noise: no padding daemon running."
258     fi
259     rm -f "$NOISE_PID"

```

```

260 }
261
262 cmd_status() {
263     echo "NOISE_ENABLED = $NOISE_ENABLED"
264     if daemon_running; then
265         echo "padding      : RUNNING (pid $(cat "$NOISE_PID"))"
266         echo "target       : ${PAD_TARGET}:${PAD_PORT}"
267         echo "mode         : ${PAD_MODE}$( [ "$PAD_MODE" = topup ] && echo " (fills to target on ${PAD_IFACE"
268         }; padding backs off as real traffic rises)"
269         echo "rate / packet : ${PAD_RATE} kbps${PAD_PCT:+ (=${PAD_PCT}% of ${PAD_LINK_MBPS} Mbps)} / ${PAD_BYTES} B, +/-${PAD_JITTER} ms jitter"
270     else
271         echo "padding      : stopped"
272     fi
273 }
274
275 # launchd entry point (see launchd/com.nullexit.noise.plist). Governed by the
276 # SAME NOISE_ENABLED switch NOT a new option. When disabled it exits 0 cleanly
277 # so KeepAlive never tight-loops; when enabled it runs the padding loop in the
278 # foreground so launchd supervises it directly. This is why persistence adds no
279 # extra knob: loading the plist once means "auto-start whenever NOISE_ENABLED=true".
280 cmd_boot() {
281     if [ "$NOISE_ENABLED" != "true" ]; then
282         log "boot: NOISE_ENABLED != true idle (padding not started)"
283         exit 0
284     fi
285     log "boot: NOISE_ENABLED=true starting padding loop under launchd"
286     pad_loop
287 }
288
289 case "${1:-}" in
290     verify)    verify_random "${2:-1048576}" loud ;;
291     start)     cmd_start ;;
292     pad)       pad_loop ;;
293     __boot)    cmd_boot ;;
294     stop)      cmd_stop ;;
295     status)    cmd_status ;;
296     --help|-h|"") sed -n '2,59p' "$0" ;;
297     *) echo "Unknown argument: $1 (try --help)" >&2; exit 2 ;;
298 esac

```

45 scripts/pf.conf

```

1 # scripts/pf.conf - Packet Filter configuration for nullexit killswitch
2 # Dynamically parameterized by pfctl macro override
3
4 # ext_if must be passed via command line, e.g. -D ext_if=en0
5 # If not passed, default to en0
6 ext_if = "en0"
7 ts_if = "utun+"
8
9 # Tables populated dynamically via pfctl
10 table <vpn_endpoints> persist
11 table <derp_relays> persist
12 # FaceTime split-route (opt-in): Apple media/relay /16s that egress DIRECT via the
13 # physical gateway. EMPTY unless FACETIME_SPLIT_ROUTE=true, so the kill-switch
14 # posture is unchanged by default. Populated by add_apple_split_routes().
15 table <apple_direct> persist
16
17 # Core Rules
18 set skip on lo0 # Ensure local loopback is never blocked
19 # TCP MSS Clamping: prevents fragmentation on host TCP flows (like Tailscale control plane)
20 # by keeping the host in sync with the container's double-tunnel (Tailscale+WARP) overhead.
21 # NOTE: this 1120 integer is rewritten at kill-switch enable time by scripts/common.sh's
22 # enable_killswitch() from GATEWAY_MSS in .env change GATEWAY_MSS in .env, not this line.
23 scrub out all max-mss 1120
24 pass quick on $ts_if all # Allow all Tailscale traffic (essential to prevent SSH lockout)
25
26 # Default deny outbound WAN traffic
27 block out on $ext_if all
28
29 # Block inbound Screen Sharing (VNC) and SSH from local Wi-Fi, forcing them to use Tailscale
30 block in on $ext_if proto tcp from any to any port { 22, 5900 }
31
32 # Exceptions for tunnel handshakes and coordination
33 pass out quick on $ext_if proto udp to <vpn_endpoints> port 2408

```

```

34 pass out quick on $ext_if proto udp to <derp_relays>
35 pass out quick on $ext_if proto tcp to 192.200.0.0/24 port 443
36
37 # FaceTime split-route (opt-in, table empty unless FACETIME_SPLIT_ROUTE=true):
38 # let real-time media/relay traffic to Apple's /16s leave DIRECT (real IP) so
39 # FaceTime can hole-punch instead of dying behind WARP's symmetric NAT.
40 pass out quick on $ext_if proto { tcp udp } to <apple_direct>
41
42 # Allow local LAN traffic (AirDrop, local Tailscale P2P, etc)
43 pass out quick on $ext_if proto tcp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
44 pass out quick on $ext_if proto udp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
45 pass out quick on $ext_if proto icmp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
46
47 # Allow DHCP outbound (client port 68 to server port 67) so we can obtain/renew IP address
48 pass out quick on $ext_if proto udp from any port 68 to any port 67

```

46 scripts/pihole-mode.sh

```

1  #!/bin/bash
2  # scripts/pihole-mode.sh  LAN Pi-hole mode (opt-in)
3  #
4  # nullexit is already a Pi-hole for MESH devices: AdGuard Home sinkholes DNS for
5  # everything that routes through the gateway over Tailscale. This mode extends
6  # that to NON-Tailscale devices on the same LAN a smart TV, a guest laptop, an
7  # IoT gadget the classic "point your DNS at this box" Pi-hole. It binds a DNS
8  # forwarder on the LAN interface that relays to AdGuard, so any device that sets
9  # its DNS server to this Mac's LAN IP gets the same ad/tracker/threatfiltering.
10 #
11 # Mechanism: reuse the tested dns-proxy.py, but LISTEN_ADDR = the LAN IP instead
12 # of loopback, forwarding to AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (the Docker
13 # port mapping). No kill-switch changes are needed pf.conf has no inbound-53
14 # block, and DNS responses to RFC1918 LAN clients are already permitted.
15 #
16 # NETWORK REALITY: on an AP-isolated network (WPA2-Enterprise / hotel), other
17 # devices CANNOT reach this Mac at all the same isolation that forces phones
18 # onto DERP. This mode only does anything useful on a TRUSTED, non-isolated
19 # network (home / your own AP). It is report-clear here but won't serve remote
20 # clients until you're on such a network. It also filters DNS only it does NOT
21 # route those devices' traffic through WARP.
22 #
23 # Usage:
24 #   bash scripts/pihole-mode.sh enable      # start the LAN DNS forwarder (needs PIHOLE_LAN_MODE=true)
25 #   bash scripts/pihole-mode.sh disable    # stop it
26 #   bash scripts/pihole-mode.sh status     # running? listen addr, AdGuard reachability
27 #   bash scripts/pihole-mode.sh test       # prove it resolves AND sinkholes, locally
28 #
29 # .env:
30 #   PIHOLE_LAN_MODE      master on/off (default false)
31 #   PIHOLE_LAN_LISTEN    LAN address to bind (blank = this Mac's en0 IP)
32 #   DNS_PROXY_PORT       AdGuard's host-published DNS port (default 5354)
33
34 set -uo pipefail
35
36 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
37 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
38 cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
39 source "$SCRIPT_DIR/common.sh"
40 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
41
42 DNS_PROXY="$SCRIPT_DIR/dns-proxy.py"
43 PID_FILE="/tmp/nullexit-pihole.pid"
44 LOG_FILE="$PROJECT_ROOT/output.log"
45
46 PIHOLE_ON=$(read_env_var "PIHOLE_LAN_MODE" | tr '[:upper:]' '[:lower:]'); [ -z "$PIHOLE_ON" ] &&
47   PIHOLE_ON="false"
48
49 # AdGuard's DNS host port is PROCEDURAL (randomized per boot see Procedural
50 # Ports), so discover it; never assume 5354. Docker is the source of truth, then
51 # the ports file toggle.sh writes, then the compose default.
52 resolve_adguard_port() {
53   local p
54   p=$(docker compose port warp 5335 2>/dev/null | grep -oE '[0-9]+$' | head -1)
55   [ -n "$p" ] && { echo "$p"; return; }
56   p=$(grep -oE 'DNS_PROXY_PORT=[0-9]+' /tmp/nullexit-ports.env 2>/dev/null | grep -oE '[0-9]+$' | head
57     -1)
58   [ -n "$p" ] && { echo "$p"; return; }

```

```

57     echo 5354
58 }
59 ADGUARD_PORT=$(resolve_adguard_port)
60
61 lan_ip() {
62     local a; a=$(read_env_var "PIHOLE_LAN_LISTEN")
63     [ -n "$a" ] && { echo "$a"; return; }
64     ipconfig getifaddr en0 2>/dev/null || echo ""
65 }
66
67 log() { echo "[$(date '+%Y-%m-%d %H:%M:%S')] [pihole-lan] $" >> "$LOG_FILE" 2>/dev/null || true; }
68
69 running() {
70     [ -f "$PID_FILE" ] || return 1
71     local p; p=$(cat "$PID_FILE" 2>/dev/null)
72     [ -n "$p" ] && kill -0 "$p" 2>/dev/null
73 }
74
75 adguard_reachable() {
76     # Probe over TCP: Colima user-mode net forwards UDP-to-container unreliably (the
77     # very reason dns-proxy.py relays over TCP), so a UDP probe false-negatives even
78     # when AdGuard is healthy. TCP is exactly the path the forwarder uses.
79     dig +tcp +short +tries=1 +time=2 -p "$ADGUARD_PORT" @127.0.0.1 example.com >/dev/null 2>&1
80 }
81
82 cmd_enable() {
83     if [ "$PIHOLE_ON" != "true" ]; then
84         echo "pihole-lan: disabled (PIHOLE_LAN_MODE != true in .env) refusing to serve LAN DNS." >&2
85         exit 3
86     fi
87     if running; then echo "pihole-lan: already running (pid $(cat "$PID_FILE"))."; exit 0; fi
88     local addr; addr=$(lan_ip)
89     if [ -z "$addr" ]; then echo "pihole-lan: could not determine a LAN IP (en0 down?)." >&2; exit 1; fi
90     if ! adguard_reachable; then
91         echo "pihole-lan: AdGuard not reachable at 127.0.0.1:$ADGUARD_PORT is the gateway up?" >&2
92         exit 1
93     fi
94     echo "pihole-lan: starting LAN DNS forwarder on ${addr}:53 AdGuard 127.0.0.1:${ADGUARD_PORT}"
95     # :53 needs root. Capture the child PID from inside the privileged shell.
96     sudo -n bash -c "LISTEN_ADDR='${addr}' LISTEN_PORT=53 TARGET_HOST=127.0.0.1 TARGET_PORT='${ADGUARD_PORT' \
97     nohup '$(command -v python3)' '$DNS_PROXY' >>'$LOG_FILE' 2>&1 & echo \${!} > '$PID_FILE'" || {
98         echo "pihole-lan: failed to start (need passwordless sudo to bind :53)." >&2; exit 1; }
99     sleep 0.4
100    if running; then
101        log "enabled on ${addr}:53 127.0.0.1:${ADGUARD_PORT}"
102        echo "pihole-lan: running (pid $(cat "$PID_FILE")). Point a LAN device's DNS at ${addr} to filter it.
103    "
104    else
105        echo "pihole-lan: listener did not stay up check $LOG_FILE (port 53 already bound?)." >&2; exit 1
106    fi
107 }
108
109 cmd_disable() {
110     if running; then
111         local p; p=$(cat "$PID_FILE")
112         sudo -n kill "$p" 2>/dev/null || kill "$p" 2>/dev/null
113         log "disabled (pid $p)"; echo "pihole-lan: stopped (pid $p)."
114     else
115         echo "pihole-lan: not running."
116     fi
117     sudo -n rm -f "$PID_FILE" 2>/dev/null || rm -f "$PID_FILE" 2>/dev/null || true
118 }
119
120 cmd_status() {
121     echo "PIHOLE_LAN_MODE = $PIHOLE_ON"
122     echo "AdGuard (127.0.0.1:$ADGUARD_PORT): $(adguard_reachable && echo reachable || echo UNREACHABLE)"
123     if running; then
124         echo "forwarder : RUNNING (pid $(cat "$PID_FILE")) on $(lan_ip):53"
125     else
126         echo "forwarder : stopped"
127     fi
128 }
129
130 cmd_test() {
131     local addr; addr=$(lan_ip)
132     if ! running; then echo "pihole-lan: not running 'enable' first." >&2; exit 1; fi
133     echo "Testing LAN Pi-hole at ${addr}:53 (querying locally)"
134     local good bad
135     good=$(dig +short +tries=1 +time=3 @"$addr" example.com 2>/dev/null | head -1)
136     bad=$(dig +short +tries=1 +time=3 @"$addr" doubleclick.net 2>/dev/null | head -1)
137     echo " resolve example.com ${good:-<none>}"

```

```

137 echo " filter doubleclick.net ${bad:-<blocked/empty>}"
138 if [ -n "$good" ] && { [ -z "$bad" ] || [ "$bad" = "0.0.0.0" ]; }; then
139     echo " RESULT: PASS resolves clean domains and sinkholes blocked ones."
140 else
141     echo " RESULT: check clean resolve=${good:+yes}, sinkhole=${bad:-empty}. (If both empty, AdGuard/
gateway may be down.)"
142 fi
143 }
144
145 case "${1:-}" in
146     enable) cmd_enable ;;
147     disable) cmd_disable ;;
148     status) cmd_status ;;
149     test) cmd_test ;;
150     --help|-h|"") sed -n '2,40p' "$0" ;;
151     *) echo "Unknown argument: $1 (try --help)" >&2; exit 2 ;;
152 esac

```

47 scripts/profiles.sh

```

1 #!/bin/bash
2 # scripts/profiles.sh named config profiles that tame the .env flag sprawl.
3 #
4 # CONFIG-ONLY BY DESIGN: this tool ONLY reads/writes .env. It NEVER touches the
5 # live gateway no routing, no PF, no containers, no restart so it cannot break
6 # your running network. Changes take effect only on your NEXT ./toggle.sh --restart.
7 #
8 # A profile is a coherent, known-good bundle of the intent-level switches, so you
9 # pick ONE intent instead of reasoning about ~7 interacting flags. KILL_SWITCH is
10 # forced true in every profile the fail-closed invariant is never negotiable.
11 #
12 # Usage:
13 #   bash scripts/profiles.sh show           # current config, grouped + explained (read-only)
14 #   bash scripts/profiles.sh list           # available profiles and what each sets (read-only)
15 #   bash scripts/profiles.sh preview <name> # exactly what 'apply' would change (read-only)
16 #   bash scripts/profiles.sh apply <name>   # write the profile to .env (backs up; never restarts)
17 #
18 # Profiles: campus (AP-isolated enterprise, conservative) home (trusted, fast)
19 #            hardened (max privacy: cover traffic + Tor bridges, no real-IP exposure)
20 #
21 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
22 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
23 cd "$PROJECT_ROOT" || exit 1
24 case "${1:-}" in --help|-h) sed -n '2,25p' "$0"; exit 0 ;; esac
25
26 python3 - "$@" <<'PY'
27 import sys, os, shutil, time
28
29 ENV = os.environ.get("PROFILES_ENV_FILE", ".env")
30
31 # Intent-level switches each profile manages. KILL_SWITCH is a forced invariant.
32 PROFILES = {
33     "campus": { "KILL_SWITCH": "true", "TAILSCALE_ALLOW_LAN_P2P": "false", "FACETIME_SPLIT_ROUTE": "false",
34                 "PIHOLE_LAN_MODE": "false", "NOISE_ENABLED": "false", "TOR_USE_BRIDGES": "false",
35                 "WARP_FAIL_THRESHOLD": "6" },
36     "home": { "KILL_SWITCH": "true", "TAILSCALE_ALLOW_LAN_P2P": "true", "FACETIME_SPLIT_ROUTE": "true",
37               "PIHOLE_LAN_MODE": "true", "NOISE_ENABLED": "false", "TOR_USE_BRIDGES": "false",
38               "WARP_FAIL_THRESHOLD": "6" },
39     "hardened": { "KILL_SWITCH": "true", "TAILSCALE_ALLOW_LAN_P2P": "false", "FACETIME_SPLIT_ROUTE": "false",
40                  "PIHOLE_LAN_MODE": "false", "NOISE_ENABLED": "true", "NOISE_PAD_MODE": "topup",
41                  "TOR_USE_BRIDGES": "true", "WARP_FAIL_THRESHOLD": "3" },
42 }
43 DESC = {
44     "campus": "AP-isolated enterprise/campus: DERP (no SNAT poisoning), everything conservative, no real-IP
45 exposure.",
46     "home": "Trusted network: direct LAN P2P (fast), FaceTime split-route + LAN Pi-hole on, no cover-traffic
47 battery cost.",
48     "hardened": "Max privacy on a hostile net: cover traffic (topup) + Tor bridges, fail-fast WARP, never
49 expose the real IP.",
50 }
51 META = {
52     "KILL_SWITCH": "PF fail-closed kill-switch (forced true never negotiable)",
53     "TAILSCALE_ALLOW_LAN_P2P": "direct LAN P2P vs forced DERP relay",
54     "FACETIME_SPLIT_ROUTE": "route Apple /i6s direct (fixes FaceTime; exposes real IP for them)",
55     "PIHOLE_LAN_MODE": "serve AdGuard DNS to non-mesh LAN devices",
56     "NOISE_ENABLED": "cover-traffic padding",

```

```

52 "NOISE_PAD_MODE": "padding rate mode (constant / topup)",
53 "TOR_USE_BRIDGES": "Tor via obfs4 bridges only",
54 "WARP_FAIL_THRESHOLD": "consecutive warp=off polls before auto-shutdown",
55 }
56
57 def read_env():
58     if not os.path.exists(ENV):
59         print(f"profiles: {ENV} not found", file=sys.stderr); sys.exit(1)
60     return open(ENV, encoding="utf-8").read().splitlines()
61
62 def get_val(lines, key):
63     for l in lines:
64         s = l.strip()
65         if s.startswith(key + "="):
66             return s.split("=", 1)[1]
67     return None
68
69 def which_profile(lines):
70     """Name the profile that matches the current .env, or None."""
71     for name, flags in PROFILES.items():
72         if all(get_val(lines, k) == v for k, v in flags.items()):
73             return name
74     return None
75
76 def cmd_show():
77     lines = read_env()
78     active = which_profile(lines)
79     print(f"\n=== nullexit config (current){'   profile: ' + active if active else '   custom (no exact'
80           'profile match)'} ===\n")
81     for k, d in META.items():
82         v = get_val(lines, k)
83         print(f"   {k:26} = {str(v):9} {d}")
84     print("\n Read-only. 'list' to see profiles, 'apply <name>' for a known-good bundle.")
85
86 def cmd_list():
87     lines = read_env()
88     print()
89     for name, flags in PROFILES.items():
90         print(f"   {name} {DESC[name]}")
91         for k, v in flags.items():
92             cur = get_val(lines, k)
93             print(f"       {k:26} {v}{'' if cur == v else '   (now: ' + str(cur) + ')'}")
94     print()
95
96 def diff_for(name):
97     lines = read_env()
98     return [(k, get_val(lines, k), v) for k, v in PROFILES[name].items() if get_val(lines, k) != v]
99
100 def cmd_preview(name):
101     if name not in PROFILES:
102         print(f"unknown profile '{name}' (try: {' , '.join(PROFILES)}", file=sys.stderr); sys.exit(2)
103     ch = diff_for(name)
104     print(f"\nprofile '{name}' {DESC[name]}\n")
105     if not ch:
106         print(" already applied no changes."); return
107     print(" would change:")
108     for k, cur, v in ch:
109         print(f"   {k:26} {cur} {v}")
110     print("\n Preview only nothing written, network untouched.")
111
112 def set_key(lines, key, val):
113     for i, l in enumerate(lines):
114         if l.strip().startswith(key + "="):
115             lead = l[:len(l) - len(l.strip())]
116             lines[i] = f"{lead}{key}={val}"
117             return
118     lines.append(f"{key}={val}")
119
120 def cmd_apply(name):
121     if name not in PROFILES:
122         print(f"unknown profile '{name}' (try: {' , '.join(PROFILES)}", file=sys.stderr); sys.exit(2)
123     ch = diff_for(name)
124     if not ch:
125         print(f"profile '{name}' already applied no changes."); return
126     lines = read_env()
127     bak = f"{ENV}.bak.{time.strftime('%Y%m%d%H%M%S')}"
128     shutil.copy2(ENV, bak)
129     for k, v in PROFILES[name].items():
130         set_key(lines, k, v)
131     with open(ENV, "w", encoding="utf-8") as f:
132         f.write("\n".join(lines) + "\n")

```



```

132     print(f"\napplied '{name}' to {ENV} (backup: {bak})\n")
133     for k, cur, v in ch:
134         print(f"        {k:26} {cur} {v}")
135     print("\n        .env only your LIVE gateway is UNCHANGED. Apply with:  ./toggle.sh --restart")
136
137 cmd = sys.argv[1] if len(sys.argv) > 1 else "show"
138 arg = sys.argv[2] if len(sys.argv) > 2 else ""
139 {"show": cmd_show, "list": cmd_list,
140  "preview": (lambda: cmd_preview(arg)), "apply": (lambda: cmd_apply(arg))}.get(cmd, cmd_show)()
141 PY

```

48 scripts/reinstall-host-tailscale.sh

```

1  #!/usr/bin/env bash
2  #
3  # scripts/reinstall-host-tailscale.sh
4  #
5  # Complete uninstall + reinstall of the host-side Homebrew Tailscale.
6  #
7  # Use this when the brew LaunchAgent is wedged, Login Items entries are
8  # duplicated, "Tailscale is stopped" persists despite
9  # 'brew services restart tailscale', or you simply want a clean baseline
10 # before re-engaging the gateway as exit node.
11 #
12 # Idempotent. Safe to re-run. Use --dry to preview every step with zero
13 # changes. Use --yes to skip the per-step confirmation prompts.
14 #
15 # This script ONLY touches the HOST-side Tailscale (the one installed by
16 # Homebrew 'brew install tailscale' and managed by launchd as
17 # homebrew.mxcl.tailscale). The nullexit gateway container's tailscaled
18 # (which lives inside the Colima VM at 100.100.21.8 and offers the exit
19 # node to this host) is unaffected.
20
21 set -u
22
23 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
24
25 DRY=0
26 YES=0
27
28 for arg in "$@"; do
29     case "$arg" in
30         --dry) DRY=1 ;;
31         --yes) YES=1 ;;
32         --help|-h) # NOTE: here doc is UNQUOTED so $$SCRIPT_DIR expands in the body. If you
33                   # add new $-style placeholders here, they'll expand the same way.
34                   cat <<USAGE
35 Usage: bash "$SCRIPT_DIR/reinstall-host-tailscale.sh" [--dry] [--yes]
36
37 --dry  Print every step, make zero changes. Use to audit before
38        committing.
39 --yes  Auto-confirm the prompts at every destructive step. Useful in
40        CI / fully-scripted recovery; do not use until you've read
41        the script.
42
43 After this script completes cleanly, you still must (MANUALLY):
44
45 1. Open System Settings Privacy & Security Local Network and
46    enable Tailscale. macOS will prompt automatically the first time
47    the fresh install of 'tailscale' is invoked.
48
49 2. Re-engage the tailnet + exit node:
50    sudo tailscale up --ssh=true --accept-dns=false \
51      --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
52      --exit-node-allow-lan-access=true
53
54 3. Re-run ./toggle.sh from the project root.
55
56 4. Final verification:
57    curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)=\'
58    expected: warp=on, ip=<Cloudflare>, colo=YYZ/YUL
59 USAGE
60     exit 0
61 ;;
62 *)
63     echo "Unknown flag: $arg" >&2

```

```

64         exit 2
65     ;;
66 esac
67 done
68
69 # ----- helpers -----
70
71 confirm() {
72     if [ "$YES" = 1 ] || [ "$DRY" = 1 ]; then
73         return 0
74     fi
75     printf "        %s [y/N] " "$1"
76     read -r ans
77     case "$ans" in
78         y|Y|yes|YES) return 0 ;;
79         *)
80             echo "        aborted by user"
81             exit 1
82             ;;
83     esac
84 }
85
86 header() {
87     printf '\n %s\n' "$1"
88 }
89
90 # with_timeout <seconds> <cmd...>
91 # Run a command in the background and kill it if it exceeds <seconds>.
92 # Returns the command's exit code if it finished in time, 124 (matching
93 # GNU 'timeout' convention) if it timed out. macOS doesn't ship GNU
94 # coreutils 'timeout', so this is the bash-3.2-friendly replacement.
95 # stdout + stderr are discarded (we only care about exit code); redirect
96 # or pipe externally if you need the output.
97 with_timeout() {
98     local timeout_s="${1:-30}"
99     shift
100     "$@" >/dev/null 2>&1 &
101     local cmdpid=$!
102     local elapsed=0
103     while kill -0 "$cmdpid" 2>/dev/null; do
104         if [ "$elapsed" -ge "$timeout_s" ]; then
105             kill -9 "$cmdpid" 2>/dev/null
106             wait "$cmdpid" 2>/dev/null
107             return 124
108         fi
109         sleep 1
110         elapsed=$((elapsed+1))
111     done
112     wait "$cmdpid" 2>/dev/null
113     return $?
114 }
115
116 # ----- 0. Pre-flight -----
117 header "Pre-flight current host tailscale state"
118 echo "    brew: $(command -v brew >/dev/null 2>&1 && brew --version | head -1 || echo '
119     not on PATH')"
119 echo "    tailscaled running: $(pgrep -lf tailscaled 2>/dev/null | wc -l | tr -d ' ') process(es)"
120 echo "    brew service tailscale: $(brew services list 2>/dev/null | awk ' $1 == \"tailscale\" {print $2, $3}
121     | tr '\n' ' ' | sed 's/ $//')'"
121 echo "    LaunchAgent plists (under ~/Library/LaunchAgents/):"
122 ls -l ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/~/ /' || echo "    (none)"
123
124 # Detect Tailscale's macOS Network Extension (residual from prior App Store
125 # Tailscale.app installs). Lives under /Library/Apple/SystemExtensions/ and
126 # is managed by 'systemextensionsctl' NOT launchd. If loaded, brew
127 # tailscaled will NOT engage the Network Extension framework slot because
128 # the SE still has it claimed. Symptom: 'sudo tailscale status' returns
129 # "Tailscale is stopped" even when the daemon is alive and listening.
130 se_found=0
131 if command -v systemextensionsctl >/dev/null 2>&1; then
132     # Grab any line matching tailscale and network-extension
133     se_line=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale\..*network-extension' |
134         head -n 1)
135     if [ -n "$se_line" ]; then
136         se_found=1
137         # Extract the teamID (10-char uppercase/numeric word)
138         se_teamID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^[A-Z0-9]{10}$' | head -n 1)
139         # Extract the bundleID (starts with io.tailscale or com.tailscale)
140         se_bundleID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^(io|com)\.tailscale\.' | head -n 1 | sed
141             's/(.*)/')

```

```

141 # If teamID is not found, fallback to '-'
142 if [ -z "$se_teamID" ]; then
143     se_teamID="-"
144 fi
145 fi
146 fi
147
148 if [ "$se_found" = 1 ] && [ -n "$se_bundleID" ]; then
149     echo "    Tailscale Network Extension System Extension is loaded (residual from App Store Tailscale.
150         app install)"
151     echo "    Team ID: $se_teamID, Bundle ID: $se_bundleID"
152     echo "    brew tailscaled cannot engage the NE slot while this SE has it claimed."
153     if [ "$DRY" = 1 ]; then
154         echo "    [dry] (would prompt) sudo systemextensionsctl uninstall $se_teamID $se_bundleID"
155     else
156         confirm "Run 'sudo systemextensionsctl uninstall $se_teamID $se_bundleID'? (frees the NE slot for
157             brew tailscale; non-fatal if it errors)"
158         sudo systemextensionsctl uninstall "$se_teamID" "$se_bundleID" 2>&1 \
159             || echo "    ! uninstall returned non-zero (may require reboot, non-fatal)"
160     fi
161 else
162     echo "    no Tailscale System Extension loaded"
163 fi
164
165 # Pattern used by the Step-1 bootout loops, in both dry + real branches.
166 # Widened from a bare /tailscale/ so the same loops also pick up any
167 # nullexit-labeled LaunchAgents (e.g., com.nullexit.wake-recovery, the
168 # caffeinate + post-wake recovery hook). The wake-recovery LaunchAgent
169 # is re-installed by ./toggle.sh / setup.sh on re-engage, so wiping it
170 # here is non-fatal and is part of restoring a clean baseline.
171 TAILSCALE_LABEL_RE='tailscale'
172 NULLEXIT_LABEL_RE='nullexit'
173
174 # ----- 1. brew services stop -----
175 header "Step 1 Stop the brew-managed service + unload LaunchAgents"
176 # Detect whether the brew-managed service is registered as $USER (gui/$UID
177 # domain) or as root (system domain). NOPASSWD sudoers can leave a stale
178 # root-owned registration after a prior 'sudo brew services ...' round;
179 # 'brew services stop' without sudo refuses to touch a root-owned plist
180 # (Error: Service 'tailscale' is started as 'root').
181 brew_status_line="$(brew services list 2>/dev/null | awk '$1=="tailscale"')"
182 brew_status_user="$(echo "$brew_status_line" | awk '{print $3}')"
183 if [ -n "$brew_status_line" ]; then
184     if [ "$DRY" = 1 ]; then
185         echo "    [dry] brew services stop tailscale (user=$brew_status_user)"
186     else
187         if [ "$brew_status_user" = "root" ]; then
188             confirm "Run 'sudo brew services stop tailscale' (service is registered as root)?"
189             sudo brew services stop tailscale 2>&1 || echo "    (sudo brew services stop returned non-zero non
190                 -fatal)"
191         else
192             echo "    brew services stop tailscale (user=$brew_status_user, no sudo needed)"
193             brew services stop tailscale 2>&1 || echo "    (brew reported it wasn't actually running)"
194         fi
195     fi
196 fi
197 echo "    unload any tailscale-related LaunchAgents (prevents launchd from respawning after kill)"
198 if [ "$DRY" = 1 ]; then
199     echo "    [dry] for each 'tailscale|nullexit'-matched label in gui/$UID domain:"
200     while IFS= read -r label; do
201         [ -z "$label" ] && continue
202         echo "    [dry] launchctl bootout gui/$UID/$label"
203     done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR
204         > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
205     echo "    [dry] for each 'tailscale|nullexit'-matched label in system domain (root-owned brew plist
206         ):"
207     while IFS= read -r label; do
208         [ -z "$label" ] && continue
209         echo "    [dry] (would prompt) sudo launchctl bootout system/$label"
210     done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$
211         $NULLEXIT_LABEL_RE" 'NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
212     if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
213         echo "    [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macos"
214     fi
215     if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
216         echo "    [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macsys"
217     fi
218 else
219     # Bootout every gui-domain LaunchAgent whose label matches 'tailscale'
220     # OR 'nullexit'. The nullexit arm catches com.nullexit.wake-recovery,
221     # which holds caffeinate/watcher.sh alive across a daemon reset not
222     # desired while the script is establishing a fresh baseline. Awake-

```

```

216 # recovery is re-installed by ./toggle.sh / setup.sh on re-engage, so
217 # wiping here is non-fatal. The pattern stays narrower than a bare
218 # /tail/ or /null/ so unrelated Apple services are not unloaded.
219 while IFS= read -r label; do
220     [ -z "$label" ] && continue
221     echo "      launchctl bootout gui/$UID/$label"
222     launchctl bootout "gui/$UID/$label" 2>/dev/null || true
223 done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR
    > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
224 # Bootout every system-domain LaunchDaemon/Agent whose label matches
225 # 'tailscale' OR 'nullexit'. This catches the root-owned brew plist left
226 # behind by prior 'sudo brew services ...' invocations that was the
227 # case that left tailscaled PID 47447 57734 respawning between sudo
228 # pkill runs. Defensive nullexit catch in case the user previously
229 # 'sudo ln -s' d the wake-recovery plist into /Library/LaunchDaemons.
230 while IFS= read -r label; do
231     [ -z "$label" ] && continue
232     confirm "Run 'sudo launchctl bootout system/$label'?"
233     sudo launchctl bootout "system/$label" 2>/dev/null || true
234 done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx=
    "$NULLEXIT_LABEL_RE" 'NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
235 # Legacy Tailscale.app system-domain entries. Each one is checked and
236 # confirmed separately so we don't sudo-promp or sudo-bootout for nothing.
237 if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
238     confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macos'?"
239     sudo launchctl bootout system/com.tailscale.ipn.macos 2>/dev/null || true
240 fi
241 if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
242     confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macsys'?"
243     sudo launchctl bootout system/com.tailscale.ipn.macsys 2>/dev/null || true
244 fi
245 fi
246 echo "      waiting up to 10s for tailscaled to exit"
247 for i in 1 2 3 4 5 6 7 8 9 10; do
248     if [ "$DRY" = 1 ]; then
249         echo "      [dry] sleep 1; loop check would terminate here"
250         break
251     fi
252     sleep 1
253     if ! pgrep -lf tailscaled >/dev/null 2>&1; then
254         echo "      tailscaled exited in ${i}s"
255         break
256     fi
257     [ "$i" = 10 ] && echo "      tailscaled still alive after 10s Step 2 will deal with it"
258 done
259 else
260     echo "      brew reports no tailscale service loaded"
261 fi
262
263 # ----- 2. Kill any orphan tailscaled -----
264 header "Step 2 Kill any orphan tailscaled process(es) (may escalate to sudo)"
265 if ! pgrep -lf tailscaled >/dev/null 2>&1; then
266     echo "      no orphans"
267 else
268     echo "      Stray process(es):"
269     pgrep -lf tailscaled | sed 's/^/ /'
270     echo "      no-sudo: pkill tailscaled"
271     if [ "$DRY" = 1 ]; then
272         echo "      [dry] pkill tailscaled on failure: sudo pkill -9 tailscaled"
273     else
274         if pkill tailscaled 2>/dev/null; then
275             echo "      no-sudo pkill succeeded"
276         else
277             echo "      ! no-sudo pkill insufficient; cannot escalate yet"
278         fi
279         sleep 1
280     fi
281 fi
282 # Recheck; if still alive, offer sudo escalation
283 if { [ "$DRY" = 0 ] && pgrep -lf tailscaled >/dev/null 2>&1; } || [ "$DRY" = 1 ]; then
284     if [ "$DRY" = 1 ]; then
285         echo "      [dry] sudo pkill -9 tailscaled (escalation would fire here)"
286     else
287         confirm "Run 'sudo pkill -9 tailscaled' to escalate?"
288         sudo pkill -9 tailscaled 2>&1
289         sleep 1
290     fi
291 fi
292
293 if [ "$DRY" = 0 ]; then
294     if pgrep -lf tailscaled >/dev/null 2>&1; then

```

```

295     echo "        Failed to kill stragglers   bailing out"
296     pgrep -lf tailscaled | sed 's/^/      /'
297     exit 3
298 fi
299     echo "        all tailscaled processes gone"
300 fi
301 fi
302
303 # ----- 3. brew uninstall -----
304 header "Step 3  brew uninstall tailscale"
305 if ! brew list tailscale >/dev/null 2>&1; then
306     echo "        not currently installed via brew  skip"
307 else
308     # Detect whether the keg directory is owner=$USER or owner=root. If the
309     # keg is root-owned (typically a residual of some prior 'sudo brew ...'
310     # cycle that did perms fixups), 'brew uninstall' without sudo refuses
311     # to remove the files and prints: 'Error: Could not remove tailscale
312     # keg! Do so manually: sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>'.
313     # Detect that and escalate to a sudoed uninstall, with a manual rm -rf
314     # fallback if 'sudo brew uninstall' itself errors out (e.g., partial
315     # state where the keg dir is gone but brew still thinks it's installed).
316     # Detect keg state. Treat '(dir absent)' AND '(dir present + empty)' as the
317     # same '<unknown>' bucket so both escalate to sudo. The empty-parent case
318     # is what you get after a manual 'sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>'
319     # from inside: brew's parent dir is left present-but-empty and brew still
320     # thinks the package is installed.
321     keg_dir="/opt/homebrew/Cellar/tailscale"
322     if [ -d "$keg_dir" ] && [ -n "$(ls -A "$keg_dir" 2>/dev/null)" ]; then
323         keg_owner="$(stat -f '%Su' "$keg_dir" 2>/dev/null || echo '<unknown>')"
324     else
325         keg_owner="<unknown>"
326     fi
327     if [ "$DRY" = 1 ]; then
328         echo "        [dry] brew uninstall tailscale (keg_owner=$keg_owner)"
329     else
330         # 'root' AND '<unknown>' both require sudo escalation. The '<unknown>'
331         # case is the partial-state path: keg dir was already removed (manually
332         # or by a prior failed uninstall) but brew's internal state still says
333         # installed. Without this branch the script would fall through to a
334         # non-sudo brew uninstall that fails again with the same 'Could not
335         # remove tailscale keg!' error.
336         if [ "$keg_owner" = "root" ] || [ "$keg_owner" = "<unknown>" ]; then
337             confirm "Run 'sudo brew uninstall tailscale' (keg=$keg_owner)?"
338             sudo brew uninstall tailscale 2>&1 || {
339                 echo "        ! sudo brew uninstall failed; offering manual rm -rf fallback"
340                 # Guard the glob so an empty/absent dir doesn't literal-expand the
341                 # '*' and confuse the user with a sudo error about a non-existent path.
342                 if [ -d /opt/homebrew/Cellar/tailscale ] && [ -n "$(ls -A /opt/homebrew/Cellar/tailscale 2>/dev/
null)" ]; then
343                     confirm "Run 'sudo rm -rf /opt/homebrew/Cellar/tailscale/*'?"
344                     sudo rm -rf /opt/homebrew/Cellar/tailscale/* 2>&1
345                 else
346                     echo "        keg dir already empty or absent  nothing more to rm"
347                 fi
348             }
349         else
350             confirm "Run 'brew uninstall tailscale' (removes LaunchAgent plist + linked symlinks)?"
351             brew uninstall tailscale 2>&1
352         fi
353     fi
354 fi
355
356 # ----- 4. Wipe state directories -----
357 header "Step 4  Wipe stale state"
358 # User-owned
359 declare -a USER_PATHS
360 USER_PATHS=(
361     "$HOME/.local/share/tailscale"
362     "$HOME/.local/run/tailscaled.socket"
363     "$HOME/.local/run/tailscaled.state"
364     "$HOME/.local/state/tailscale"
365     "$HOME/Library/Application Support/Tailscale"
366     "$HOME/Library/Caches/com.tailscale.ipn.macos"
367     "$HOME/Library/Caches/Tailscale"
368     "$HOME/Library/Logs/Tailscale"
369 )
370 echo "        User-owned paths (no sudo):"
371 for p in "${USER_PATHS[@]"; do
372     if [ -e "$p" ]; then
373         if [ "$DRY" = 1 ]; then
374             echo "        [dry] rm -rf '$p'"

```

```

375     else
376         rm -rf "$p"
377         echo "      rm -rf '$p'"
378     fi
379 fi
380 done
381
382 # System
383 declare -a SUDO_PATHS
384 SUDO_PATHS=(
385     "/var/lib/tailscale"
386     "/var/cache/tailscale"
387 )
388 echo "      System paths under /var (will prompt for sudo if any are present):"
389 eligible_sudo=0
390 for p in "${SUDO_PATHS[@]"; do
391     if [ -e "$p" ]; then
392         eligible_sudo=1
393         if [ "$DRY" = 1 ]; then
394             echo "      [dry] sudo rm -rf '$p'"
395         else
396             if confirm "Run 'sudo rm -rf $p'?" ; then
397                 sudo rm -rf "$p" && echo "      rm -rf '$p'"
398             fi
399         fi
400     fi
401 done
402 [ "$eligible_sudo" = 0 ] && echo "      (none of /var/lib/tailscale, /var/cache/tailscale are present)"
403
404 # ----- 5. Wipe stale LaunchAgent plists -----
405 header "Step 5 Wipe stale LaunchAgent plist files"
406 declare -a PLIST_FILES
407 PLIST_FILES=(
408     "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
409     "$HOME/Library/LaunchAgents/com.tailscale.ipn.macos.plist"
410     "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist.bak"
411     "$HOME/Library/LaunchAgents/com.nullexit.wake-recovery.plist"
412 )
413 for f in "${PLIST_FILES[@]"; do
414     if [ -e "$f" ]; then
415         if [ "$DRY" = 1 ]; then
416             echo "      [dry] rm -f '$f'"
417         else
418             rm -f "$f"
419             echo "      rm -f '$f'"
420         fi
421     fi
422 done
423 echo "      Remaining tailscale|nullexit plists (should be empty):"
424 ls -l ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/~/ /' || echo "      (none)"
425
426 # ----- 5.5. Drain Background-Items (Login Items) -----
427 header "Step 5.5 Drain Background-Items (Login Items)"
428 # macOS 13+ keeps a parallel "Background Items" database alongside
429 # classic LaunchAgents. Even when we wipe ~/Library/LaunchAgents/homebrew.mxcl.tailscale.plist
430 # in Step 5, the .btm file at
431 # ~/Library/Application Support/com.apple.backgroundtaskmanagementagent/backgrounditems.btm
432 # can still hold a Background-Items entry pointing at the (now-dead) plist
433 # path. Next login: macOS sees the entry, tries to bootstrap, finds the
434 # plist missing, and either silently regenerates it via brew's post-install
435 # hook (causing tailscaled respawn at next login) or refuses outright.
436 # 'sfltool reset-login-items' is Apple's canonical API for wiping the
437 # entire .btm file. It DOES wipe all Login Items the user re-adds what
438 # they want via System Settings General Login Items. Acceptable here
439 # because the goal of this script is "fresh baseline for Tailscale".
440 BTM="$HOME/Library/Application Support/com.apple.backgroundtaskmanagementagent/backgrounditems.btm"
441 if [ -f "$BTM" ]; then
442     if [ "$DRY" = 1 ]; then
443         echo "      [dry] found $BTM; would prompt sfltool reset-login-items"
444     else
445         confirm "Run 'sfltool reset-login-items' (wipes ALL Login Items; rebuild via System Settings General Login Items)?"
446         if sfltool reset-login-items 2>&1; then
447             echo "      Login Items drained"
448         else
449             echo "      ! sfltool reset-login-items failed (non-fatal; do manually: System Settings General Login Items)"
450         fi
451     fi
452 else
453     echo "      no backgrounditems.btm present"

```

```

454 fi
455
456 # ----- 6. brew install -----
457 header "Step 6 brew install tailscale (fresh)"
458 if brew list tailscale >/dev/null 2>&1; then
459     echo "        already installed skip"
460 else
461     confirm "Run 'brew install tailscale'?"
462     if [ "$DRY" = 1 ]; then
463         echo "        [dry] brew install tailscale"
464     else
465         brew install tailscale 2>&1
466     fi
467 fi
468
469 # ----- 6.5. Strip quarantine xattr from keg -----
470 header "Step 6.5 Strip quarantine xattr from tailscale install"
471 # brew's freshly poured bottles don't carry com.apple.quarantine (the
472 # bottles are content-addressed and Homebrew-verified), but the upstream
473 # installer (Tailscale.app from App Store) and any manual download do.
474 # Strip the xattr from the entire keg as a defensive safety net so the
475 # freshly-installed tailscaled binary and any helper binaries are not
476 # Gatekeeper-blocked on first launch. Idempotent no-op if no quarantine
477 # xattr is present.
478 quarantine_count="$(xattr -lr /opt/homebrew/Cellar/tailscale 2>/dev/null | grep -c '^[^ ]* com.apple.
479 quarantine' || echo 0)"
480 if [ "${quarantine_count:-0}" -gt 0 ]; then
481     if [ "$DRY" = 1 ]; then
482         echo "        [dry] found $quarantine_count file(s) with com.apple.quarantine; would prompt xattr -dr"
483     else
484         confirm "Run 'xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale'? (strips quarantine
485         xattr from $quarantine_count file(s); non-fatal if it errors)"
486         if xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale 2>&1; then
487             echo "        quarantine xattr stripped"
488         else
489             echo "        ! xattr -dr failed (non-fatal; open /opt/homebrew/Cellar/tailscale in Finder, right-click
490             tailscaled Open Anyway)"
491         fi
492     fi
493 else
494     echo "        no com.apple.quarantine xattr on tailscale install"
495 fi
496
497 # ----- 7. brew services start (NO sudo!) -----
498 header "Step 7 Start the service as $USER (NO sudo, with multi-tier self-healing)"
499 if ! command -v tailscale >/dev/null 2>&1; then
500     echo "        tailscale not on PATH after install brew install errored; aborting"
501     exit 4
502 fi
503
504 # Aliveness fan-out: ANY one of these is taken as proof of life. None of
505 # them alone is reliable pgrep misses tailscaled if it does setproctitle
506 # or fork-detach; launchctl's reported PID can be a shim that exited
507 # within 1s; tailscale-cli hangs while macOS Local-Network permission is
508 # unresolved. Combined, they cover each other's blind spots.
509
510 tailscaled_api_up() {
511     # Strategy A: tailscale CLI can talk to the local API. Canonical proof.
512     # Bound at 8s tailscale status itself returns quickly on success OR
513     # failure, but if the post-install state is wedged it can hang.
514     with_timeout 8 /opt/homebrew/bin/tailscale status >/dev/null 2>&1
515 }
516
517 tailscaled_proc_up() {
518     # Strategy B: any process has 'tailscaled' substring in argv.
519     pgrep -lf tailscaled >/dev/null 2>&1
520 }
521
522 tailscaled_lc_up() {
523     # Strategy C: launchctl-managed service has a non-dash, non-zero PID.
524     local lc_pid
525     lc_pid="$(launchctl list 2>/dev/null | awk ' $3 == "homebrew.mxcl.tailscale" {print $1; exit; }' )"
526     if [ -n "$lc_pid" ] && [ "$lc_pid" != "-" ] && [ "$lc_pid" -gt 0 ] 2>/dev/null; then
527         return 0
528     fi
529     return 1
530 }
531
532 tailscaled_is_alive() {
533     tailscaled_api_up && return 0
534     tailscaled_proc_up && return 0
535 }

```



```

532     tailscaled_lc_up    && return 0
533     return 1
534 }
535
536 confirm "Run 'brew services start tailscale'?"
537 if [ "$DRY" = 1 ]; then
538     echo "    [dry] brew services start tailscale"
539     echo "    [dry] (then poll up to 30s for tailscaled_is_alive)"
540     echo "    [dry] (recovery tier 1: launchctl kickstart)"
541     echo "    [dry] (recovery tier 2: launchctl bootout + bootstrap of explicit plist)"
542     echo "    [dry] (recovery tier 3: brew services restart (stop+start cycle))"
543     echo "    [dry] (recovery tier 4: direct /opt/homebrew/bin/tailscaled spawn (bypasses launchd))"
544 else
545     brew services start tailscale 2>&1
546     # Primary poll tailscaled_is_alive covers all 3 detection strategies.
547     echo "    polling for tailscaled to come up (up to ~30s)"
548     tailscaled_up=0
549     for i in 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15; do
550         if tailscaled_is_alive; then
551             tailscaled_up=1
552             echo "    tailscaled alive after ${i}*2s"
553             break
554         fi
555         sleep 2
556     done
557
558     # Recovery tier 1: kickstart the launchd-managed service. This is the
559     # cheapest and most-aligned-with-launchd fix when brew has loaded the
560     # LaunchAgent but tailscaled didn't fork. Kickstart tells launchd
561     # "re-spawn if not running" without forcing a teardown.
562     if [ "$tailscaled_up" = 0 ]; then
563         echo "    ! Recovery tier 1: launchctl kickstart"
564         if launchctl kickstart -kp "gui/$UID/homebrew.mxcl.tailscale" 2>&1; then
565             for i in 1 2 3 4 5 6 7 8 9 10; do
566                 if tailscaled_is_alive; then
567                     tailscaled_up=1
568                     echo "    tailscaled alive after kickstart (${i}*2s)"
569                     break
570                 fi
571                 sleep 2
572             done
573         else
574             echo "    ! kickstart returned non-zero (LaunchAgent may not be loaded)"
575         fi
576     fi
577
578     # Recovery tier 2: explicit plist bootout + bootstrap. Forces a fresh
579     # launchd registration cycle for the brew-managed service, which fixes
580     # the wider class of "brew says started but daemon never spawned"
581     # wedge states that kickstart alone can't dig out of.
582     if [ "$tailscaled_up" = 0 ]; then
583         echo "    ! Recovery tier 2: launchctl bootout + bootstrap of explicit plist"
584         launchctl bootout "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null || true
585         sleep 1
586         # Brew installs the canonical plist template at $HOMEBREW_PREFIX/Library/.
587         # On Apple Silicon that's /opt/homebrew/Library/LaunchAgents/. On Intel
588         # it's /usr/local/Library/LaunchAgents/. Detect which and pick the
589         # right one so this script works across both architectures.
590         if [ -f /opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
591             plist_path="/opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
592         elif [ -f /usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
593             plist_path="/usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
594         else
595             plist_path=""
596         fi
597         if [ -n "$plist_path" ]; then
598             if launchctl bootstrap "gui/$UID" "$plist_path" 2>&1; then
599                 for i in 1 2 3 4 5 6 7 8 9 10; do
600                     if tailscaled_is_alive; then
601                         tailscaled_up=1
602                         echo "    tailscaled alive after bootstrap (${i}*2s)"
603                         break
604                     fi
605                     sleep 2
606                 done
607             else
608                 echo "    ! explicit bootstrap returned non-zero"
609             fi
610         else
611             echo "    ! could not locate brew plist template; skipping bootstrap tier"
612         fi

```

```

613 fi
614
615 # Recovery tier 3: brew services restart (forces a stop + start cycle,
616 # which unsticks LaunchAgents with stale 'started' state from prior
617 # half-cycles). This is the canonical brew-doc fix for 'services don't
618 # actually start even though brew says started'.
619 if [ "$tailscaled_up" = 0 ]; then
620     echo "    ! Recovery tier 3: brew services restart (forces stop+start cycle)"
621     if brew services restart tailscale 2>&1; then
622         for i in 1 2 3 4 5 6 7 8 9 10; do
623             if tailscaled_is_alive; then
624                 tailscaled_up=1
625                 echo "        tailscaled alive after brew restart (${i}*2s)"
626                 break
627             fi
628             sleep 2
629         done
630     else
631         echo "        ! brew services restart returned non-zero"
632     fi
633 fi
634
635 # Recovery tier 4: last-resort direct spawn of the daemon binary,
636 # bypassing launchd entirely. If tailscaled crashes here too, we'll
637 # capture its actual stderr in /tmp for the user to read (NOPASSWD
638 # sudoers often have local-network permission denied the captured
639 # stderr makes it obvious what to fix in System Settings).
640 if [ "$tailscaled_up" = 0 ]; then
641     echo "    ! Recovery tier 4: direct spawn (bypasses launchd; tail stderr to /tmp)"
642     /opt/homebrew/bin/tailscaled >/tmp/nullexit-tailscaled-spawn.log 2>&1 &
643     spawned_pid=$!
644     disown "$spawned_pid" 2>/dev/null || true
645     sleep 3
646     if tailscaled_is_alive; then
647         tailscaled_up=1
648         echo "        tailscaled alive via direct spawn (pid=$spawned_pid)"
649         echo "        (NOTE: not under launchd supervision a reboot, re-run of"
650         echo "        this script, or ./toggle.sh is required to recreate."
651         echo "        Last 5 lines of stderr captured to /tmp/nullexit-tailscaled-spawn.log:)"
652         tail -n 5 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/' '/' || true
653     fi
654 fi
655
656 if [ "$tailscaled_up" = 0 ]; then
657     echo "    HARD FAIL: tailscaled refuses to stay alive across every recovery tier"
658     echo "    launchctl-managed state (if available):"
659     launchctl print "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null \
660     | grep -E 'state =|last exit code =|runs =|path =' \
661     | sed 's/^/' '/' \
662     || echo "        (no launchctl state available for homebrew.mxcl.tailscale)"
663     echo "    Last 10 lines of /tmp/nullexit-tailscaled-spawn.log (if any):"
664     tail -n 10 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/' '/' || echo "    (no log)"
665     echo "    Most likely remaining cause: macOS Local Network permission denied."
666     echo "    System Settings Privacy & Security Local Network toggle Tailscale ON."
667     exit 7
668 fi
669 fi
670 sleep 3
671
672 # ----- 8. Verify -----
673 header "Step 8 Verify fresh install"
674 echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{{print $2, $3}}' | tr '\n' ' ' | sed 's/ $//')"
```

```

691 echo "      tailscale status:      (binary missing or errored)"
692 fi
693
694 if { [ "$DRY" = 0 ] && /opt/homebrew/bin/tailscale status 2>/dev/null | grep -q '^Tailscale is stopped';
    }; then
695     echo
696     echo "      NOTICE: 'tailscale status' still reports stopped after a fresh install."
697     echo "      Two equally likely causes:"
698     echo "      (a) macOS hasn't yet shown the first-run 'Local Network' prompt for"
699     echo "      Tailscale open System Settings Privacy & Security Local"
700     echo "      Network. Toggle Tailscale ON."
701     echo "      (b) macOS denied Local Network permission at some prior point."
702     echo "      Either way: System Settings Privacy & Security Local Network."
703     echo "      If the entry isn't present, invoking any tailscale command will"
704     echo "      re-trigger the prompt automatically."
705 fi
706
707 # ----- 9. Done -----
708 header "Done"
709 if [ "$DRY" = 1 ]; then
710     echo "      (dry-mode finished; no follow-up commands were issued.)"
711     exit 0
712 fi
713
714 cat <<'DONE'
715
716 Manual follow-ups (all of these; the script stops short of running sudo
717 tailscale up so you can verify the brew install is actually live first):
718
719 1. System Settings Privacy & Security Local Network Tailscale ON
720    macOS will normally pop the prompt the first time you run 'tailscale'.
721    If it didn't, trigger it via the menu bar check.
722
723 2. Re-engage the tailnet + exit node:
724    sudo tailscale up --ssh=true --accept-dns=false \
725      --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
726      --exit-node-allow-lan-access=true
727
728 3. Re-run: ./toggle.sh
729    (re-installs the com.nullexit.wake-recovery LaunchAgent wiped
730    in Step 5 to clear the slate and re-engages the gateway)
731
732 4. Final verification:
733    curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)=\'
734    expect: warp=on, ip=<Cloudflare IP>, colo=YYZ/YUL
735
736 Container-side (gateway / 100.100.21.8) Tailscale is unaffected by this
737 script. Other tailnet devices (S24, iPhone, Windows) are unaffected.
738
739 DONE

```

49 scripts/routing-fix.sh

```

1 #!/bin/sh
2 # routing-fix: Maintains routing for SOCKS5 proxy traffic through WARP (tun0)
3 # and FORWARD rules for Tailscale exit-node return traffic.
4 #
5 # Runs inside the warp container's network namespace. This script ensures:
6 # 1. Table 200 has default via tun0 (for SOCKS5 proxy traffic from 172.18.0.2)
7 # with the WARP endpoint exception via eth0 (prevents tunnel loop)
8 # 2. The CGNAT ip rule (100.64.0.0/10 table 52 tailscale routes) is maintained for Tailscale
9 # 3. FORWARD RELATED,ESTABLISHED rule allows return traffic for exit-node
10 # forwarded connections (S24 tailscale0 eth0 host internet)
11 #
12 # CRITICAL: Does NOT touch the main table's default route. Gluetun handles that
13 # internally via its own routing rules and WireGuard fwmark (0xca6c table 51820).
14 # Overriding the main table default would create a loop: encrypted WireGuard
15 # packets exiting tun0 would match default dev tun0 and re-enter the tunnel.
16
17 set -e
18
19 trap 'echo "routing-fix: Caught signal, exiting..."; exit 0' TERM INT QUIT
20
21 sleep 15 &
22 wait $! || true
23

```

```

24 echo "routing-fix: Setting up routing for SOCKS5 proxy through WARP (tun0)..."
25
26 # Dynamically detect Docker subnet from eth0
27 DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')
28 DOCKER_GW=$(ip route show default | head -1 | awk '{print $3}')
29 WARP_ENDPOINT="${WARP_ENDPOINT}"
30 TAILSCALE_ALLOW_LAN_P2P="${TAILSCALE_ALLOW_LAN_P2P:-false}"
31 IP_BLOCKLIST_FILE="/userfilters/ip_blocklist.ipset"
32 LAST_IP_MTIME=""
33 echo "routing-fix: Docker network: ${DOCKER_NET}, gateway: ${DOCKER_GW}"
34
35 # Table 200: Used by ip rule 100 (from 172.18.0.2) for SOCKS5 proxy traffic
36 # WARP endpoint goes via eth0 to avoid tunnel loop, everything else via tun0
37 ip route add "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || \
38 ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
39 ip route replace default dev tun0 table 200 2>/dev/null || true
40
41 # Main table: Just ensure the Docker subnet route exists via eth0
42 # (gluetun owns the default route here do NOT touch it)
43 ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true
44
45 # FORWARD RELATED, ESTABLISHED: Allow return traffic for exit-node connections.
46 # The container has BOTH iptables (nftables backend) and iptables-legacy
47 # loaded. Gluetun's post-rules.txt uses the nftables backend, while Tailscale
48 # uses the legacy backend. Packets go through BOTH stacks, so we add the rule
49 # to both to ensure it survives regardless of which backend has policy DROP.
50 #
51 # Packet flow: S24 outgoing (tailscale0->eth0) ACCEPTed by first rule.
52 # Return traffic (eth0->eth0 via table 52 to host) needs RELATED, ESTABLISHED
53 # to match the conntrack entry created by the outgoing flow.
54 add_fwd_related_established() {
55     # nftables backend (used by gluetun's post-rules.txt)
56     if ! iptables -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
57         iptables -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
58     fi
59     # legacy backend (used by Tailscale's ts-forward)
60     if command -v iptables-legacy >/dev/null 2>&1; then
61         if ! iptables-legacy -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
62             iptables-legacy -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
63         fi
64     fi
65 }
66
67 add_fwd_related_established
68
69 # Policy Routing for Tailscale P2P:
70 # Tailscale sends its encrypted mesh UDP packets from port 41642.
71 # If these go out tun0 (WARP), Cloudflare's strict NAT drops the returning
72 # hole-punch packets, causing Tailscale to fail P2P and fall back to DERP relays (high latency).
73 # We mark packets originating from port 41642 and force them out the main table (eth0).
74 #
75 # LAN P2P is controlled by two sources (in priority order):
76 # 1. /app/.lan_p2p_detected written by watcher.sh on macOS on every network change
77 #    using system_profiler SPAirPortDataType (802.1x detection) + AP isolation probe. Overrides .env.
78 # 2. TAILSCALE_ALLOW_LAN_P2P env var (from .env) static fallback.
79 # When disabled, we explicitly DROP local RFC1918-destined Tailscale UDP so the
80 # phone cleanly falls back to DERP rather than getting poisoned by macOS gvproxy SNAT.
81 add_tailscale_p2p_bypass() {
82     # Dynamic override from watcher.sh (network auto-detection) takes priority over .env
83     local allow_p2p="${TAILSCALE_ALLOW_LAN_P2P:-false}"
84     if [ -f "/app/.lan_p2p_detected" ]; then
85         allow_p2p=$(cat "/app/.lan_p2p_detected" 2>/dev/null | tr -d '[:space:]')
86     fi
87
88     # Always allow P2P directly to the Docker host gateway (the Mac running this container).
89     # The gateway container is always co-located with the host Mac on the same machine,
90     # so direct Tailscale P2P via the Docker bridge is always safe and avoids DERP overhead.
91     # This RETURN rule must be inserted BEFORE the general DROP rules so it takes priority.
92     local default_gw
93     default_gw=$(ip route show default | awk '{print $3}')
94     if [ -n "$default_gw" ]; then
95         if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null;
96         then
97             iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null ||
98             true
99         fi
100     fi
101
102     # Also explicitly whitelist all physical IPs of the host Mac (e.g. 172.17.52.223).
103     # The Tailscale control plane registers the Mac using its physical IP, so the
104     # container will attempt P2P handshakes using that physical IP instead of the

```

```

103 # Docker bridge IP. If this physical IP falls within 172.16.0.0/12, it would be dropped.
104 if [ -f "/app/.host_ips" ]; then
105     while IFS= read -r host_ip; do
106         [ -z "$host_ip" ] && continue
107         if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null;
108         then
109             iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null ||
110             true
111         fi
112     done < "/app/.host_ips"
113 fi
114
115 if [ "${allow_p2p}" != "true" ]; then
116     # Drop Tailscale UDP packets destined for local subnets to prevent macOS gvproxy SNAT
117     # endpoint poisoning. When disabled (campus/enterprise WPA2-Enterprise with AP isolation),
118     # dropping these forces the S24 to use DERP relays reliably.
119     for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
120         if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; then
121             iptables -t mangle -A OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null || true
122         fi
123     done
124 else
125     # Remove DROP rules if they exist (switching from restricted to allowed)
126     for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
127         while iptables -t mangle -D OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; do ;;
128         done
129     done
130 fi
131
132 # Bypass STUN and P2P packets to public IPs (DERP relays and remote peers) out eth0
133 if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null; then
134     iptables -t mangle -A OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null || true
135 fi
136 if ! ip rule show | grep -q 'fwmark 0x8888'; then
137     ip rule add fwmark 0x8888 lookup 254 pref 98 2>/dev/null || true
138 fi
139 }
140
141 add_tailscale_p2p_bypass
142
143 add_onion_transparent_proxy() {
144     for ipt in iptables iptables-legacy; do
145         command -v "$ipt" >/dev/null 2>&1 || continue
146         if ! $ipt -t nat -C PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null; then
147             $ipt -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null || true
148         fi
149         if ! $ipt -C FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null; then
150             $ipt -I FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null || true
151         fi
152     done
153 }
154
155 add_tailscale_p2p_bypass
156 add_onion_transparent_proxy
157
158 create_log_drop_chain() {
159     local ipt="$1"
160     local chain="$2"
161     local prefix="$3"
162
163     if ! $ipt -N "$chain" 2>/dev/null; then
164         $ipt -F "$chain" 2>/dev/null || true
165     fi
166     $ipt -A "$chain" -j LOG --log-prefix "$prefix"
167     $ipt -A "$chain" -j DROP 2>/dev/null || true
168 }
169
170 add_country_block() {
171     # To add a country to the blocklist, simply define the 'BLOCKED_COUNTRIES' variable in .env with 2-
172     # letter ISO country codes.
173     # You can find the country code by looking at the filename on http://www.ipdeny.com/ipblocks/data/countries/
174     # For example, to block China (cn) and Russia (ru), set: BLOCKED_COUNTRIES="cn ru" in your .env file
175     for zone in $BLOCKED_COUNTRIES; do
176         set_name="block_${zone}"
177         zone_upper=$(echo "$zone" | tr 'a-z' 'A-Z')
178         chain_name="LOG_DROP_${zone_upper}"
179         log_prefix="IP_BLOCK_${zone_upper}: "
180
181         # Create the ipset if it doesn't exist
182         if ! ipset list "$set_name" >/dev/null 2>&1; then

```

```

179 ipset create "$set_name" hash:net 2>/dev/null || true
180 echo "routing-fix: Downloading ${zone_upper} IPs for blocklist..."
181 curl -sS "http://www.ipdeny.com/ipblocks/data/countries/${zone}.zone" | grep -v '^#' | while read -
r subnet; do
182     [ -n "$subnet" ] && ipset add "$set_name" "$subnet" 2>/dev/null || true
183 done
184 fi
185
186 # Apply LOG_DROP rule to both backends
187 for ipt in iptables iptables-legacy; do
188     command -v "$ipt" >/dev/null 2>&1 || continue
189
190     create_log_drop_chain "$ipt" "$chain_name" "$log_prefix"
191
192     # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
193     while $ipt -D FORWARD -m set --match-set "$set_name" dst -j DROP 2>/dev/null; do ;; done
194
195     if ! $ipt -C FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null; then
196         $ipt -I FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null || true
197     fi
198 done
199 }
200
201 add_ip_blocklist() {
202     # Only reload when the compiled file has actually changed (mtime check).
203     # This prevents a pointless ipset restore on every 30-second loop tick.
204     if [ ! -f "$IP_BLOCKLIST_FILE" ]; then
205         return 0
206     fi
207
208     CURRENT_MTIME=$(stat -c %Y "$IP_BLOCKLIST_FILE" 2>/dev/null || echo "0")
209     if [ "$CURRENT_MTIME" != "$LAST_IP_MTIME" ]; then
210         echo "routing-fix: IP blocklist changed, reloading..."
211         # ipset restore handles the atomic swap internally:
212         # create_new populate swap with live destroy_new
213         if ipset restore < "$IP_BLOCKLIST_FILE" 2>/dev/null; then
214             LAST_IP_MTIME="$CURRENT_MTIME"
215             echo "routing-fix: IP blocklist loaded ($(ipset list block_malicious | grep -c '[0-9]') entries)."
216         else
217             echo "routing-fix: Warning: ipset restore failed. Retrying next cycle."
218             return 1
219         fi
220     fi
221
222     # Apply FORWARD rules in BOTH iptables backends.
223     for ipt in iptables iptables-legacy; do
224         command -v "$ipt" >/dev/null 2>&1 || continue
225
226         create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_DST "IP_BLOCK_MALICIOUS_DST: "
227         create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_SRC "IP_BLOCK_MALICIOUS_SRC: "
228
229         # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
230         while $ipt -D FORWARD -m set --match-set block_malicious dst -j DROP 2>/dev/null; do ;; done
231         while $ipt -D FORWARD -m set --match-set block_malicious src -j DROP 2>/dev/null; do ;; done
232
233         if ! $ipt -C FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null;
234         then
235             $ipt -I FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null ||
236             true
237         fi
238
239         if ! $ipt -C FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null;
240         then
241             $ipt -I FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null ||
242             true
243         fi
244     done
245 }
246
247 # Start background block logger (DNS + IP)
248 if ! pgrep -f "nullexit-logger" >/dev/null; then
249     echo "routing-fix: Starting background block logger..."
250     nullexit-logger &
251 fi
252
253 add_country_block
254 add_ip_blocklist
255
256 echo 'routing-fix: Routes applied.'
```

```

255 UNHEALTHY_COUNT=0
256
257 # Re-assert loop (every 30 seconds)
258 while true; do
259     # Table 200 routes
260     ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
261
262     # Tunnel health check
263     if ! curl -m 3 -sS --interface tun0 https://cloudflare.com/cdn-cgi/trace >/dev/null 2>&1; then
264         UNHEALTHY_COUNT=$((UNHEALTHY_COUNT + 1))
265     else
266         UNHEALTHY_COUNT=0
267         if [ -f "/app/TUNNEL_FAILED_CLOSED.marker" ]; then
268             rm -f "/app/TUNNEL_FAILED_CLOSED.marker"
269         fi
270     fi
271
272     if [ "$UNHEALTHY_COUNT" -ge 3 ]; then
273         echo "routing-fix: Tunnel health check failed $UNHEALTHY_COUNT times. Failing closed."
274         ip route del default dev tun0 table 200 2>/dev/null || true
275         ip route replace blackhole default table 200 2>/dev/null || true
276         touch "/app/TUNNEL_FAILED_CLOSED.marker"
277     else
278         ip route replace default dev tun0 table 200 2>/dev/null || true
279     fi
280
281     # Main table: keep Docker subnet route
282     ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true
283
284     # CGNAT ip rule for Tailscale
285     if ! ip rule show | grep -q '100.64.0.0/10'; then
286         ip rule add to 100.64.0.0/10 lookup 52 pref 99 2>/dev/null || true
287         echo 'routing-fix: CGNAT rule missing, re-injected'
288     fi
289
290     # FORWARD RELATED, ESTABLISHED: Allow return traffic in both iptables
291     # backends. Gluetun may reset the nftables ruleset on VPN reconnect,
292     # and Tailscale may reset the legacy ruleset on auth refresh.
293     add_fwd_related_established
294
295     # Ensure P2P packets bypass WARP to maintain low latency
296     add_tailscale_p2p_bypass
297
298     # Ensure transparent proxying for .onion is active
299     add_onion_transparent_proxy
300
301     # Enforce country blocklist
302     add_country_block
303
304     # Enforce IP blocklist (reloads automatically when file changes)
305     add_ip_blocklist
306
307     sleep 30 &
308     wait $! || true
309 done

```

50 scripts/rule-compiler/go.mod

```

1 module nullexit/rule-compiler
2
3 go 1.22

```

51 scripts/rule-compiler/main.go

```

1 package main
2
3 import (
4     "bufio"
5     "bytes"
6     "crypto/md5"
7     "encoding/base64"
8     "encoding/hex"

```



```

9  "fmt"
10 "io"
11 "net/http"
12 "net/netip"
13 "os"
14 "os/exec"
15 "path/filepath"
16 "regexp"
17 "sort"
18 "strings"
19 "sync"
20 "time"
21 )
22
23 const (
24     BlackListPath = "black_list.txt"
25     WhiteListPath = "white_list.txt"
26     OutputPath    = "adguard/work/userfilters/compiled_rules.txt"
27     IpOutputPath  = "adguard/work/userfilters/ip_blocklist.ipset"
28     IpCacheDir    = "adguard/work/userfilters/cache/ip"
29     CacheDir      = "adguard/work/userfilters/cache"
30 )
31
32 var CoreLists = []string{
33     "https://raw.githubusercontent.com/anudeepND/blacklist/master/adservers.txt",
34     "https://raw.githubusercontent.com/anudeepND/blacklist/master/facebook.txt",
35     "https://raw.githubusercontent.com/crazy-max/WindowsSpyBlocker/master/data/hosts/spy.txt",
36     "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.samsung.txt",
37     "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.apple.txt",
38     "https://abp.oisd.nl/basic/",
39     "https://adaway.org/hosts.txt",
40     "https://pgl.yoyo.org/adserver/serverlist.php?hostformat=adblockplus&showintro=1&mimetype=plaintext",
41     "https://raw.githubusercontent.com/nextdns/cname-cloaking-blocklist/master/domains",
42 }
43
44 var MediumAdditions = []string{
45     "https://raw.githubusercontent.com/lightswitch05/hosts/master/docs/lists/facebook-extended.txt",
46     "https://someonewhocares.org/hosts/zero/hosts",
47     "https://urlhaus.abuse.ch/downloads/hostfile/",
48 }
49
50 var HeavyAdditions = []string{
51     "https://raw.githubusercontent.com/StevenBlack/hosts/master/hosts",
52     "https://raw.githubusercontent.com/Perflyst/PiHoleBlocklist/master/SmartTV-AGH.txt",
53     "https://raw.githubusercontent.com/DandelionSprout/adfilt/master/GameConsoleAdblockList.txt",
54 }
55
56 var IpBlockLists = []string{
57     "https://feodotracker.abuse.ch/downloads/ipblocklist.txt",
58     "https://www.spamhaus.org/drop/drop.txt",
59     "https://rules.emergingthreats.net/fwrules/emerging-Block-IPs.txt",
60     "https://cinsscore.com/list/ci-badguys.txt",
61 }
62
63 func getProfiles() map[string][]string {
64     profiles := make(map[string][]string)
65
66     light := append([]string(nil), CoreLists...)
67     light = append(light, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/light.txt")
68     profiles["light"] = light
69
70     medium := append([]string(nil), CoreLists...)
71     medium = append(medium, MediumAdditions...)
72     medium = append(medium, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/multi.txt")
73     profiles["medium"] = medium
74
75     heavy := append([]string(nil), CoreLists...)
76     heavy = append(heavy, MediumAdditions...)
77     heavy = append(heavy, HeavyAdditions...)
78     heavy = append(heavy, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/pro.txt")
79     profiles["heavy"] = heavy
80
81     return profiles
82 }
83
84 func loadEnvProfile() string {
85     profile := "medium"
86
87     if bytesData, err := os.ReadFile(".env"); err == nil {
88         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))

```

```

89     for scanner.Scan() {
90         line := strings.TrimSpace(scanner.Text())
91         if line != "" && !strings.HasPrefix(line, "#") {
92             parts := strings.SplitN(line, "=", 2)
93             if len(parts) == 2 {
94                 key := strings.TrimSpace(parts[0])
95                 val := strings.TrimSpace(parts[1])
96                 if key == "GATEWAY_RULE_PROFILE" {
97                     val = strings.Trim(val, "\\")
98                     profile = strings.ToLower(val)
99                 }
100             }
101         }
102     }
103 } else if !os.IsNotExist(err) {
104     fmt.Printf("Warning: Failed to read .env file for profile configuration (%v). Using default.\n", err)
105 }
106
107 if envVal := os.Getenv("GATEWAY_RULE_PROFILE"); envVal != "" {
108     profile = strings.ToLower(envVal)
109 }
110
111 profiles := getProfiles()
112 if _, exists := profiles[profile]; !exists {
113     fmt.Printf("Warning: Profile '%s' is invalid. Falling back to 'medium'.\n", profile)
114     profile = "medium"
115 }
116 return profile
117 }
118
119 func loadDomains(filepath string) map[string]bool {
120     domains := make(map[string]bool)
121     if bytesData, err := os.ReadFile(filepath); err == nil {
122         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
123         for scanner.Scan() {
124             line := strings.TrimSpace(scanner.Text())
125             if line != "" && !strings.HasPrefix(line, "#") {
126                 domains[strings.ToLower(line)] = true
127             }
128         }
129     }
130     return domains
131 }
132
133 func parseDomainsFromContent(content string) map[string]bool {
134     domains := make(map[string]bool)
135     scanner := bufio.NewScanner(strings.NewReader(content))
136     ipRegex := regexp.MustCompile(`^\d{1,3}\.\d{1,3}\.\d{1,3}\.\d{1,3}$`)
137     agRegex := regexp.MustCompile(`^\\|\\|([a-z0-9.-]+)\\$`)
138     domainRegex := regexp.MustCompile(`^[a-z0-9.-]+$`)
139
140     for scanner.Scan() {
141         line := strings.TrimSpace(scanner.Text())
142         if line == "" || strings.HasPrefix(line, "!") || strings.HasPrefix(line, "[") {
143             continue
144         }
145         parts := strings.SplitN(line, "#", 2)
146         line = strings.TrimSpace(parts[0])
147         if line == "" {
148             continue
149         }
150
151         fields := strings.Fields(line)
152         var rule string
153         if len(fields) >= 2 {
154             if ipRegex.MatchString(fields[0]) {
155                 rule = fields[1]
156             } else {
157                 rule = fields[0]
158             }
159         } else {
160             rule = fields[0]
161         }
162
163         rule = strings.TrimSpace(strings.ToLower(rule))
164
165         var domain string
166         if m := agRegex.FindStringSubmatch(rule); m != nil {
167             domain = m[1]
168         } else if domainRegex.MatchString(rule) {
169             domain = rule

```

```

170     } else {
171         domain = rule
172     }
173
174     if domain != "" && domain != "localhost" && domain != "0.0.0.0" && domain != "127.0.0.1" && domain !=
175         "broadcasthost" {
176         domains[domain] = true
177     }
178     return domains
179 }
180
181 // WARNING: Parsing AdGuardHome.yaml using regex instead of a proper YAML parser
182 // is brittle and assumes the exact format currently emitted by AdGuard.
183 // If an AdGuard version bump updates the configuration schema format,
184 // this parser will fail silently or loudly and should be refactored to use gopkg.in/yaml.v3.
185 func getAdguardNativeLists() []string {
186     yamlPath := "adguard/conf/AdGuardHome.yaml"
187     var enabledUrls []string
188     if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
189         return enabledUrls
190     }
191
192     contentBytes, err := os.ReadFile(yamlPath)
193     if err != nil {
194         fmt.Printf("Warning: Failed to parse AdGuardHome.yaml for native lists (%v)\n", err)
195         return enabledUrls
196     }
197     content := string(contentBytes)
198
199     filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*)"?:~[a-zA-Z_]+:~\z`)
200     match := filtersRegex.FindStringSubmatch(content)
201     if len(match) > 1 {
202         filtersText := match[1]
203         blocks := strings.Split(filtersText, "\n - ")
204         urlRegex := regexp.MustCompile(`url:\s*(\S+)`)
205
206         for _, block := range blocks {
207             if strings.TrimSpace(block) == "" {
208                 continue
209             }
210             if strings.Contains(block, "enabled: true") {
211                 urlMatch := urlRegex.FindStringSubmatch(block)
212                 if len(urlMatch) > 1 {
213                     url := urlMatch[1]
214                     if !strings.Contains(url, "compiled_rules.txt") {
215                         enabledUrls = append(enabledUrls, url)
216                     }
217                 }
218             }
219         }
220     }
221     return enabledUrls
222 }
223
224 func getAdguardCompiledRulesCachePath() string {
225     yamlPath := "adguard/conf/AdGuardHome.yaml"
226     if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
227         return ""
228     }
229
230     contentBytes, err := os.ReadFile(yamlPath)
231     if err != nil {
232         fmt.Printf("Warning: Failed to resolve compiled_rules.txt cache path (%v)\n", err)
233         return ""
234     }
235     content := string(contentBytes)
236
237     filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*)"?:~[a-zA-Z_]+:~\z`)
238     match := filtersRegex.FindStringSubmatch(content)
239     if len(match) > 1 {
240         blocks := strings.Split(match[1], "\n - ")
241         idRegex := regexp.MustCompile(`id:\s*(\d+)`)
242         for _, block := range blocks {
243             if !strings.Contains(block, "compiled_rules.txt") {
244                 continue
245             }
246             idMatch := idRegex.FindStringSubmatch(block)
247             if len(idMatch) > 1 {
248                 return fmt.Sprintf("adguard/work/data/filters/%s.txt", idMatch[1])
249             }

```

```

250 }
251 }
252 return ""
253 }
254
255 func handleFetchError(urlStr string, cacheFile string, err error) map[string]bool {
256     if _, statErr := os.Stat(cacheFile); statErr == nil {
257         fmt.Printf(" -> Warning: Failed to fetch %s (%v). Falling back to expired local cache.\n", urlStr,
258             err)
259         content, readErr := os.ReadFile(cacheFile)
260         if readErr == nil {
261             domains := parseDomainsFromContent(string(content))
262             fmt.Printf(" -> Loaded from expired cache (%d domains).\n", len(domains))
263             return domains
264         }
265         fmt.Printf(" -> Warning: Failed to read expired cache (%v). Using local lists only.\n", readErr)
266     } else {
267         fmt.Printf(" -> Warning: Failed to fetch %s (%v). Using local lists only for this source.\n", urlStr,
268             err)
269     }
270     return make(map[string]bool)
271 }
272
273 func fetchRemoteDomains(urlStr string) map[string]bool {
274     os.MkdirAll(CacheDir, 0755)
275
276     hash := md5.Sum([]byte(urlStr))
277     urlHash := hex.EncodeToString(hash[:])
278     cacheFile := filepath.Join(CacheDir, fmt.Sprintf("%s.txt", urlHash))
279
280     if stat, err := os.Stat(cacheFile); err == nil {
281         fileAge := time.Since(stat.ModTime()).Seconds()
282         if fileAge < 86400 {
283             fmt.Printf("Loading remote blacklist from cache: %s\n", urlStr)
284             content, err := os.ReadFile(cacheFile)
285             if err == nil {
286                 domains := parseDomainsFromContent(string(content))
287                 fmt.Printf(" -> Loaded from cache (copy is %.1f hours old, %d domains).\n", fileAge/3600.0, len(
288                     domains))
289                 return domains
290             }
291             fmt.Printf(" -> Warning: Failed to read cache file %s (%v). Re-fetching...\n", cacheFile, err)
292         }
293     }
294
295     fmt.Printf("Fetching remote blacklist from: %s ...\n", urlStr)
296
297     req, err := http.NewRequest("GET", urlStr, nil)
298     if err != nil {
299         return handleFetchError(urlStr, cacheFile, err)
300     }
301     req.Header.Set("User-Agent", "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7)")
302
303     client := &http.Client{Timeout: 15 * time.Second}
304     resp, err := client.Do(req)
305     if err != nil {
306         return handleFetchError(urlStr, cacheFile, err)
307     }
308     defer resp.Body.Close()
309
310     if resp.StatusCode != 200 {
311         return handleFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
312     }
313
314     bodyBytes, err := io.ReadAll(resp.Body)
315     if err != nil {
316         return handleFetchError(urlStr, cacheFile, err)
317     }
318
319     content := string(bodyBytes)
320     domains := parseDomainsFromContent(content)
321
322     if len(domains) < 10 {
323         return handleFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only found %d domains.
324             Possible 404 or bad URL", len(domains)))
325     }
326
327     err = os.WriteFile(cacheFile, bodyBytes, 0644)
328     if err != nil {
329         fmt.Printf(" -> Warning: Failed to save cache file (%v)\n", err)
330     }
331 }

```

```

327
328     fmt.Printf(" -> Successfully fetched and cached %d domains.\n", len(domains))
329     return domains
330 }
331
332 func optimizeSubdomains(domains map[string]bool, listName string) map[string]bool {
333     fmt.Printf("Optimizing %s by removing redundant subdomains...\n", listName)
334     rawCount := len(domains)
335     if rawCount == 0 {
336         return make(map[string]bool)
337     }
338
339     optimized := make(map[string]bool)
340     for domain := range domains {
341         if strings.ContainsAny(domain, "|^@$/") {
342             optimized[domain] = true
343             continue
344         }
345
346         parts := strings.Split(domain, ".")
347         hasParent := false
348
349         for i := 1; i < len(parts)-1; i++ {
350             parent := strings.Join(parts[i:], ".")
351             if domains[parent] {
352                 hasParent = true
353                 break
354             }
355         }
356
357         if !hasParent {
358             optimized[domain] = true
359         }
360     }
361
362     saved := rawCount - len(optimized)
363     reduction := float64(0)
364     if rawCount > 0 {
365         reduction = (float64(saved) / float64(rawCount)) * 100
366     }
367     fmt.Printf(" -> Reduced %s from %d to %d domains (-%d / %.1f%% reduction).\n", listName, rawCount, len(
368         optimized), saved, reduction)
369     return optimized
370 }
371
372 func parseIpsFromContent(content string) map[string]bool {
373     ips := make(map[string]bool)
374     scanner := bufio.NewScanner(strings.NewReader(content))
375
376     for scanner.Scan() {
377         line := strings.TrimSpace(scanner.Text())
378         if line == "" || strings.HasPrefix(line, "#") || strings.HasPrefix(line, ";") {
379             continue
380         }
381
382         fields := strings.Fields(line)
383         if len(fields) > 0 {
384             token := strings.TrimRight(fields[0], ";,")
385             if _, err := netip.ParsePrefix(token); err == nil {
386                 ips[token] = true
387             } else if _, err := netip.ParseAddr(token); err == nil {
388                 ips[token] = true
389             }
390         }
391     }
392     return ips
393 }
394
395 func handleIpFetchError(urlStr string, cacheFile string, err error) map[string]bool {
396     if _, statErr := os.Stat(cacheFile); statErr == nil {
397         fmt.Printf(" -> Warning: fetch failed (%v). Falling back to stale cache.\n", err)
398         content, readErr := os.ReadFile(cacheFile)
399         if readErr == nil {
400             ips := parseIpsFromContent(string(content))
401             fmt.Printf(" -> %d IPs loaded from stale cache.\n", len(ips))
402             return ips
403         }
404         fmt.Printf(" -> Warning: stale cache also failed (%v).\n", readErr)
405     } else {
406         fmt.Printf(" -> Warning: %s failed (%v). No cache available. Skipping.\n", urlStr, err)
407     }
408 }

```

```

407     return make(map[string]bool)
408 }
409
410 func fetchRemoteIps(urlStr string) map[string]bool {
411     os.MkdirAll(IpCacheDir, 0755)
412
413     hash := md5.Sum([]byte(urlStr))
414     urlHash := hex.EncodeToString(hash[:])
415     cacheFile := filepath.Join(IpCacheDir, fmt.Sprintf("%s.txt", urlHash))
416
417     if stat, err := os.Stat(cacheFile); err == nil {
418         fileAge := time.Since(stat.ModTime()).Seconds()
419         if fileAge < 86400 {
420             fmt.Printf("Loading IP feed from cache: %s\n", urlStr)
421             content, err := os.ReadFile(cacheFile)
422             if err == nil {
423                 ips := parseIpsFromContent(string(content))
424                 fmt.Printf(" -> %d IPs (cache is %.1fh old).\n", len(ips), fileAge/3600.0)
425                 return ips
426             }
427             fmt.Printf(" -> Warning: cache read failed (%v). Re-fetching...\n", err)
428         }
429     }
430
431     fmt.Printf("Fetching IP feed: %s ...\n", urlStr)
432     req, err := http.NewRequest("GET", urlStr, nil)
433     if err != nil {
434         return handleIpFetchError(urlStr, cacheFile, err)
435     }
436     req.Header.Set("User-Agent", "Mozilla/5.0")
437
438     client := &http.Client{Timeout: 15 * time.Second}
439     resp, err := client.Do(req)
440     if err != nil {
441         return handleIpFetchError(urlStr, cacheFile, err)
442     }
443     defer resp.Body.Close()
444
445     if resp.StatusCode != 200 {
446         return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
447     }
448
449     bodyBytes, err := io.ReadAll(resp.Body)
450     if err != nil {
451         return handleIpFetchError(urlStr, cacheFile, err)
452     }
453
454     content := string(bodyBytes)
455     ips := parseIpsFromContent(content)
456     if len(ips) < 1 {
457         return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only %d IPs found. Possible 404", len(ips)))
458     }
459
460     err = os.WriteFile(cacheFile, bodyBytes, 0644)
461     if err != nil {
462         fmt.Printf(" -> Warning: cache write failed (%v)\n", err)
463     }
464
465     fmt.Printf(" -> %d IPs fetched and cached.\n", len(ips))
466     return ips
467 }
468
469 func compileIpBlocklist() int {
470     fmt.Println("\n IP Blocklist Compilation ")
471
472     allIps := make(map[string]bool)
473     var mu sync.Mutex
474     var wg sync.WaitGroup
475
476     maxThreads := 8
477     if len(IpBlockLists) < 8 {
478         maxThreads = len(IpBlockLists)
479     }
480     semaphore := make(chan struct{}, maxThreads)
481
482     for _, urlStr := range IpBlockLists {
483         wg.Add(1)
484         go func(u string) {
485             defer wg.Done()
486             semaphore <- struct{}{}

```

```

487     defer func() { <-semaphore }()
488
489     ips := fetchRemoteIps(u)
490     mu.Lock()
491     for ip := range ips {
492         allIps[ip] = true
493     }
494     mu.Unlock()
495 } (urlStr)
496 }
497 wg.Wait()
498
499 fmt.Printf("Total unique entries before filtering: %d\n", len(allIps))
500
501 reservedStr := []string{
502     "10.0.0.0/8",
503     "172.16.0.0/12",
504     "192.168.0.0/16",
505     "127.0.0.0/8",
506     "169.254.0.0/16",
507     "100.64.0.0/10",
508     "0.0.0.0/8",
509 }
510 var reserved []netip.Prefix
511 for _, s := range reservedStr {
512     p, _ := netip.ParsePrefix(s)
513     reserved = append(reserved, p)
514 }
515
516 cleanIps := make(map[string]bool)
517 for entry := range allIps {
518     var prefix netip.Prefix
519     p, err := netip.ParsePrefix(entry)
520     if err != nil {
521         addr, err2 := netip.ParseAddr(entry)
522         if err2 != nil {
523             continue
524         }
525         prefix = netip.PrefixFrom(addr, addr.BitLen())
526     } else {
527         prefix = p
528     }
529
530     overlaps := false
531     for _, r := range reserved {
532         if r.Overlaps(prefix) {
533             overlaps = true
534             break
535         }
536     }
537     if !overlaps {
538         cleanIps[entry] = true
539     }
540 }
541
542 removed := len(allIps) - len(cleanIps)
543 if removed > 0 {
544     fmt.Printf(" -> Removed %d private/reserved entries.\n", removed)
545 }
546 fmt.Printf(" -> Final IP blocklist: %d entries.\n", len(cleanIps))
547
548 os.MkdirAll(filepath.Dir(IpOutputPath), 0755)
549
550 f, err := os.Create(IpOutputPath)
551 if err != nil {
552     fmt.Printf("Error writing IP output file: %v\n", err)
553     return len(cleanIps)
554 }
555 defer f.Close()
556
557 f.WriteString("# nullexit Compiled IP Blocklist\n")
558 f.WriteString("# Sources: Feodo Tracker, Spamhaus DROP/EDROP, Emerging Threats, CINS\n")
559 f.WriteString(fmt.Sprintf("# Entries: %d\n\n", len(cleanIps)))
560
561 f.WriteString("create block_malicious_new hash:net maxelem 200000 -exist\n")
562
563 var sortedIps []string
564 for ip := range cleanIps {
565     sortedIps = append(sortedIps, ip)
566 }
567 sort.Strings(sortedIps)

```



```

568
569     for _, ip := range sortedIps {
570         f.WriteString(fmt.Sprintf("add block_malicious_new %s -exist\n", ip))
571     }
572
573     f.WriteString("create block_malicious hash:net maxelem 200000 -exist\n")
574     f.WriteString("swap block_malicious block_malicious_new\n")
575     f.WriteString("destroy block_malicious_new\n")
576
577     fmt.Printf("IP blocklist written to %s\n", IpOutputPath)
578     return len(cleanIps)
579 }
580
581 func main() {
582     profile := loadEnvProfile()
583     fmt.Printf("Active memory profile: '%s'\n", strings.ToUpper(profile))
584
585     localBlackList := loadDomains(BlackListPath)
586     whiteList := loadDomains(WhiteListPath)
587
588     blackList := make(map[string]bool)
589     for k := range localBlackList {
590         blackList[k] = true
591     }
592
593     profiles := getProfiles()
594     urlsToFetch := profiles[profile]
595
596     fmt.Printf("\nStarting concurrent downloads for %d lists...\n", len(urlsToFetch))
597     startTime := time.Now()
598
599     var mu sync.Mutex
600     var wg sync.WaitGroup
601
602     maxThreads := 16
603     if len(urlsToFetch) < 16 {
604         maxThreads = len(urlsToFetch)
605     }
606     semaphore := make(chan struct{}, maxThreads)
607
608     for _, urlStr := range urlsToFetch {
609         wg.Add(1)
610         go func(u string) {
611             defer wg.Done()
612             semaphore <- struct{}{}
613             defer func() { <-semaphore }()
614
615             domains := fetchRemoteDomains(u)
616             mu.Lock()
617             for d := range domains {
618                 blackList[d] = true
619             }
620             mu.Unlock()
621         }(urlStr)
622     }
623     wg.Wait()
624
625     fmt.Printf("Finished concurrent downloads in %.2f seconds using %d threads.\n", time.Since(startTime).
        Seconds(), maxThreads)
626
627     adguardNativeUrls := getAdguardNativeLists()
628     var adguardNativeDomains map[string]bool
629     if len(adguardNativeUrls) > 0 {
630         fmt.Printf("\nFetching %d AdGuard native list(s) for deduplication...\n", len(adguardNativeUrls))
631         adguardNativeDomains = make(map[string]bool)
632
633         maxNativeThreads := 4
634         if len(adguardNativeUrls) < 4 {
635             maxNativeThreads = len(adguardNativeUrls)
636         }
637         nativeSemaphore := make(chan struct{}, maxNativeThreads)
638
639         for _, urlStr := range adguardNativeUrls {
640             wg.Add(1)
641             go func(u string) {
642                 defer wg.Done()
643                 nativeSemaphore <- struct{}{}
644                 defer func() { <-nativeSemaphore }()
645
646                 domains := fetchRemoteDomains(u)
647                 mu.Lock()

```

```

648         for d := range domains {
649             adguardNativeDomains[d] = true
650         }
651         mu.Unlock()
652     }(urlStr)
653 }
654 wg.Wait()
655
656 if len(adguardNativeDomains) > 0 {
657     fmt.Println("Deduplicating compiled rules against AdGuard native lists...")
658     originalSize := len(blackList)
659     for d := range adguardNativeDomains {
660         delete(blackList, d)
661     }
662     fmt.Printf(" -> Removed %d redundant rules already covered by AdGuard.\n", originalSize-len(
        blackList))
663 }
664 }
665
666 var contradictions []string
667 for d := range whiteList {
668     if blackList[d] {
669         contradictions = append(contradictions, d)
670     }
671 }
672
673 if len(contradictions) > 0 {
674     fmt.Printf("Detected %d contradictions. Whitelist taking priority.\n", len(contradictions))
675     sort.Strings(contradictions)
676     for _, domain := range contradictions {
677         if localBlackList[domain] {
678             fmt.Printf(" -> Removed local blacklist domain: %s\n", domain)
679         }
680         delete(blackList, domain)
681     }
682 }
683
684 blackList = optimizeSubdomains(blackList, "blacklist")
685 whiteList = optimizeSubdomains(whiteList, "whitelist")
686
687 os.MkdirAll(filepath.Dir(OutputPath), 0755)
688 outFile, err := os.Create(OutputPath)
689 if err != nil {
690     fmt.Printf("Error creating output file: %v\n", err)
691     return
692 }
693 defer outFile.Close()
694
695 outFile.WriteString("! Custom Compiled Rules (Auto-Generated)\n")
696 outFile.WriteString(fmt.Sprintf("! Memory Profile: %s\n", strings.ToUpper(profile)))
697 outFile.WriteString(fmt.Sprintf("! Total Block Rules: %d\n", len(blackList)))
698 outFile.WriteString(fmt.Sprintf("! Native AdGuard Rules: %d\n", len(adguardNativeDomains)))
699 outFile.WriteString(fmt.Sprintf("! Total Whitelist Rules: %d\n", len(whiteList)))
700
701 outFile.WriteString("! --- Blacklist Rules ---\n")
702
703 var sortedBlacklist []string
704 for d := range blackList {
705     sortedBlacklist = append(sortedBlacklist, d)
706 }
707 sort.Strings(sortedBlacklist)
708
709 for _, domain := range sortedBlacklist {
710     if strings.Contains(domain, "$") {
711         parts := strings.SplitN(domain, "$", 2)
712         dom, mod := parts[0], parts[1]
713         if strings.HasPrefix(dom, "|") {
714             if strings.HasSuffix(dom, "~") {
715                 outFile.WriteString(fmt.Sprintf("%s%s\n", dom, mod))
716             } else {
717                 outFile.WriteString(fmt.Sprintf("%s~%s\n", dom, mod))
718             }
719         } else {
720             outFile.WriteString(fmt.Sprintf("||s~%s\n", dom, mod))
721         }
722     } else if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(
723     domain, "|") || strings.HasSuffix(domain, "~") {
724         outFile.WriteString(fmt.Sprintf("%s\n", domain))
725     } else {
726         outFile.WriteString(fmt.Sprintf("||s~\n", domain))

```

```

727 }
728
729 outFile.WriteString("\n! --- Whitelist Rules ---\n")
730 var sortedWhitelist []string
731 for d := range whiteList {
732     sortedWhitelist = append(sortedWhitelist, d)
733 }
734 sort.Strings(sortedWhitelist)
735
736 for _, domain := range sortedWhitelist {
737     if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|")
738     || strings.HasSuffix(domain, "^") || strings.HasPrefix(domain, "@@") {
739         rule := domain
740         if !strings.HasPrefix(domain, "@@") {
741             rule = "@@" + domain
742         }
743         outFile.WriteString(fmt.Sprintf("%s\n", rule))
744     } else {
745         outFile.WriteString(fmt.Sprintf("@@||%s^\n", domain))
746     }
747 }
748
749 fmt.Printf("\nSuccessfully compiled %d block rules and %d allow rules to %s\n", len(blackList), len(
750     whiteList), OutputPath)
751
752 cachedFilterPath := getAdguardCompiledRulesCachePath()
753 if cachedFilterPath != "" {
754     input, err := os.ReadFile(OutputPath)
755     if err == nil {
756         err = os.WriteFile(cachedFilterPath, input, 0644)
757         if err == nil {
758             fmt.Printf("Updated AdGuard filter cache: %s\n", cachedFilterPath)
759         } else {
760             fmt.Printf("Warning: Could not update AdGuard filter cache %s (%v)\n", cachedFilterPath, err)
761         }
762     } else {
763         fmt.Printf("Warning: Could not read OutputPath to copy to cache (%v)\n", err)
764     }
765 } else {
766     fmt.Println("Warning: Could not determine compiled_rules.txt cache path from AdGuardHome.yaml;
767     skipping cache update.")
768 }
769
770 adguardReachable := false
771
772 dockerComposeYamlExists := false
773 if _, err := os.Stat("docker-compose.yml"); err == nil {
774     dockerComposeYamlExists = true
775 }
776
777 dockerPath, err := exec.LookPath("docker")
778 if err == nil && dockerComposeYamlExists {
779     cmd := exec.Command(dockerPath, "compose", "ps", "--status", "running", "-q", "adguardhome")
780     out, err := cmd.Output()
781     if err == nil && strings.TrimSpace(string(out)) != "" {
782         fmt.Println("Force-recreating AdGuard Home container to drop persisted filter state...")
783         upCmd := exec.Command(dockerPath, "compose", "up", "-d", "--force-recreate", "--no-deps", "
784             adguardhome")
785         err := upCmd.Run()
786         if err == nil {
787             fmt.Println("AdGuard Home force-recreated successfully.")
788             time.Sleep(8 * time.Second)
789             adguardReachable = true
790         } else {
791             fmt.Println("Failed to restart AdGuard Home. Is it running?")
792         }
793     }
794 }
795
796 if adguardReachable {
797     agPass := "nullexit"
798     if bytesData, err := os.ReadFile(".env"); err == nil {
799         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
800         for scanner.Scan() {
801             line := scanner.Text()
802             if strings.HasPrefix(line, "ADGUARD_PASSWORD=") {
803                 parts := strings.SplitN(line, "=", 2)
804                 if len(parts) == 2 {
805                     agPass = strings.Trim(strings.TrimSpace(parts[1]), "\"'")
806                 }
807             }
808             break
809         }
810     }
811 }

```

```

804 }
805 }
806
807 creds := base64.StdEncoding.EncodeToString([]byte(fmt.Sprintf("admin:%s", agPass)))
808 req, err := http.NewRequest("POST", "http://127.0.0.1:3000/control/filtering/refresh", bytes.
809   NewBuffer([]byte("{}")))
810 if err == nil {
811     req.Header.Set("Authorization", fmt.Sprintf("Basic %s", creds))
812     req.Header.Set("Content-Type", "application/json")
813     client := &http.Client{Timeout: 15 * time.Second}
814     resp, err := client.Do(req)
815     if err == nil {
816         fmt.Printf("Triggered AdGuard filter refresh via REST API (HTTP=%d).\n", resp.StatusCode)
817         resp.Body.Close()
818     } else {
819         fmt.Printf("Warning: AdGuard filter refresh API call failed (%v). New whitelist will not load
820           until next refresh tick (~24h) or manual refresh.\n", err)
821     }
822 }
823 }
824 }
825
826 compileIpBlocklist()
827 }

```

52 scripts/setup-common.sh

```

1 #!/bin/bash
2 # 1. Docker
3 step "Checking Docker"
4
5 if ! command -v docker >> output.log 2>&1; then
6     echo ""
7     if [[ "$OSTYPE" == "darwin"* ]]; then
8         echo " Docker is not installed."
9         # Note: Homebrew aggressively rolls forward and discourages version pinning.
10        # This setup is confirmed stable on Colima v0.10.3 and Docker v29.6.1.
11        read -rp " Would you like to automatically install Colima & Docker via Homebrew? [y/N]: "
12        INSTALL_DOCKER
13        if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
14            step "Installing Colima and Docker..."
15            brew install colima docker docker-compose
16            step "Starting Colima VM..."
17            colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged
18        else
19            echo ""
20            echo " Please install Docker manually. Two options:"
21            echo " A) Colima (Recommended): brew install colima docker docker-compose && colima start --
22            memory 0.6 --vm-type vz --network-address --network-mode bridged"
23            echo " B) Docker Desktop: https://www.docker.com/products/docker-desktop/"
24            die "Install Docker, then re-run this script."
25        fi
26    else
27        echo " Docker is not installed."
28        # Note: The Docker convenience script always fetches the latest stable release.
29        # This setup is confirmed stable on Docker Engine v29.6.1.
30        read -rp " Would you like to automatically install Docker? (Requires sudo) [y/N]: "
31        INSTALL_DOCKER
32        if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
33            step "Installing Docker..."
34            curl -fsSL https://get.docker.com | sh
35            sudo usermod -aG docker "$USER"
36            echo -e "\n\033[0;32m[OK] Docker installed successfully!\033[0m"
37            echo -e "\033[0;33m[!] CRITICAL: You must log out of your SSH session and log back in for
38            Docker permissions to apply.\033[0m"
39            die "Please log out, log back in, and re-run setup.sh."
40        else
41            echo " Please install Docker manually:"
42            echo " curl -fsSL https://get.docker.com | sh"
43            echo " sudo usermod -aG docker $USER && newgrp docker"
44            die "Install Docker, then re-run this script."
45        fi
46    fi
47 fi
48
49 if ! docker info >> output.log 2>&1; then
50     echo ""

```

```

47 if [[ "$OSTYPE" == "darwin"* ]]; then
48     echo " Docker is installed but not running."
49     echo " If using Colima:          colima start --memory 0.6 --vm-type vz --network-address --
network-mode bridged"
50     echo " If using Docker Desktop: open the app."
51 else
52     echo " Docker is installed but not running."
53     echo " Run: sudo systemctl start docker"
54 fi
55 die "Start Docker, then re-run this script."
56 fi
57
58 if ! docker compose version >> output.log 2>&1; then
59     die "Docker Compose v2 not found. Update Docker or install the compose plugin:\n https://docs.docker
.com/compose/install/"
60 fi
61
62 ok "Docker $(docker --version | grep -oE '[0-9.]+' | head -1) is running."
63
64 # 1.5 Python 3
65 step "Checking Python 3"
66
67 if ! command -v python3 >> output.log 2>&1; then
68     echo ""
69     if [[ "$OSTYPE" == "darwin"* ]]; then
70         echo " Python 3 is not installed."
71         echo " macOS includes it via Xcode Command Line Tools, or you can install it via Homebrew:"
72         echo " brew install python3"
73         die "Install Python 3, then re-run this script."
74     else
75         echo " Python 3 is not installed."
76         echo " Please install it using your system's package manager. For example:"
77         echo " Debian/Ubuntu: sudo apt install python3"
78         echo " Fedora: sudo dnf install python3"
79         echo " Arch: sudo pacman -S python"
80         die "Install Python 3, then re-run this script."
81     fi
82 fi
83
84 # Note: This setup is confirmed stable on Python 3.9.6 (macOS default).
85 PYTHON_VER=$(python3 --version 2>&1 | head -n 1)
86 ok "$PYTHON_VER is installed."
87
88 # 2. Tailscale (host)
89 step "Checking Tailscale (host)"
90 # The host needs its own Tailscale installation separate from the Docker
91 # container so Toggle-Gateway can run 'tailscale down' before stopping
92 # containers and 'tailscale up' after starting them. Without this, the
93 # exit-node deadlock that setup.sh is designed to prevent can still occur.
94
95 # Resolve the tailscale binary: CLI (brew/package) or bundled macOS app
96 TAILSCALE_BIN=""
97 if command -v tailscale >> output.log 2>&1; then
98     TAILSCALE_BIN="tailscale"
99 elif [[ -x "/Applications/Tailscale.app/Contents/MacOS/Tailscale" ]]; then
100     TAILSCALE_BIN="/Applications/Tailscale.app/Contents/MacOS/Tailscale"
101 fi
102
103 RECOMMENDED_TS_VERSION="1.98.5"
104
105 if [[ -z "$TAILSCALE_BIN" ]]; then
106     warn "Tailscale not found on host installing..."
107
108     if [[ "$OSTYPE" == "darwin"* ]]; then
109         if command -v brew >> output.log 2>&1; then
110             step "Installing Tailscale CLI formula via Homebrew (standalone, no .app GUI)..."
111             # Note: We install Tailscale via Homebrew, targeting the stable version (recommended: 1.98.5)
112             brew install tailscale -q
113             TAILSCALE_BIN="tailscale"
114
115             step "Starting tailscaled as a per-user LaunchAgent (auto-starts at login)..."
116             if brew services start tailscale 2>&1 | grep -qE 'Successfully|started'; then
117                 ok "tailscaled LaunchAgent registered."
118             else
119                 warn "Could not auto-start tailscaled. Run 'brew services start tailscale' manually."
120                 warn "For a system-wide LaunchDaemon that runs before login, run instead:"
121                 warn " sudo brew services start tailscale"
122             fi
123             echo ""
124         else

```

```

125     die "Homebrew not found. Install Tailscale manually (recommended version: ${
RECOMMENDED_TS_VERSION}): \n  https://tailscale.com/download/mac \n  Then re-run this script."
126     fi
127
128     elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
129         curl -fsSL https://tailscale.com/install.sh | sh
130         sudo systemctl enable --now tailscaled 2>> output.log || true
131         TAILSCALE_BIN="tailscale"
132
133     else
134         die "Unsupported OS. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION})
: \n  https://tailscale.com/download \n  Then re-run this script."
135     fi
136
137     TAILSCALE_VER=$((${TAILSCALE_BIN} version 2>> output.log | head -n 1)
138     if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
139         warn "Tailscale installed, but version ($TAILSCALE_VER) differs from recommended version (
$RECOMMENDED_TS_VERSION)."
140     else
141         ok "Tailscale installed (version $TAILSCALE_VER)."
142     fi
143 else
144     TAILSCALE_VER=$((${TAILSCALE_BIN} version 2>> output.log | head -n 1)
145     if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
146         warn "Tailscale is already installed ($TAILSCALE_VER), but recommended version is
$RECOMMENDED_TS_VERSION for host/container consistency."
147     else
148         ok "Tailscale is already installed at recommended version ($TAILSCALE_VER)."
149     fi
150 fi
151
152 # Authenticate the host machine if it isn't already connected.
153 # This is an interactive step it opens a browser login URL.
154 # We do it now, before containers start, so Toggle-Gateway can
155 # safely run 'tailscale down/up' from day one.
156 if ! ${TAILSCALE_BIN} status >> output.log 2>&1; then
157     echo ""
158     echo "  Your host machine needs to join your Tailscale network."
159     echo "  A browser window or login URL will appear authenticate, then"
160     echo "  return here. The script will continue automatically."
161     echo ""
162     if [[ "$OSTYPE" == "linux-gnu"* ]]; then
163         sudo ${TAILSCALE_BIN} up
164     else
165         ${TAILSCALE_BIN} up
166     fi
167     ok "Host machine connected to Tailscale."
168 else
169     ok "Host machine is already connected to Tailscale."
170 fi
171
172 # 3. wgcf
173 step "Checking wgcf (Cloudflare WARP key tool)"
174
175 if ! command -v wgcf >> output.log 2>&1; then
176     warn "wgcf not found installing..."
177
178     if [[ "$OSTYPE" == "darwin"* ]]; then
179         command -v brew >> output.log 2>&1 || die "Homebrew not found. \n Install wgcf manually: https://
github.com/ViRb3/wgcf/releases"
180         brew install wgcf -q
181
182     elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
183         ARCH=$(uname -m)
184         case $ARCH in
185             x86_64) WGCF_ARCH="amd64" ;;
186             aarch64) WGCF_ARCH="arm64" ;;
187             armv7l) WGCF_ARCH="armv7" ;;
188             *) die "Unsupported architecture: $ARCH. \n Install wgcf manually: https://github.com/ViRb3/
wgcf/releases" ;;
189         esac
190
191         echo "  Fetching latest release..."
192         WGCF_VERSION=$(curl -sf https://api.github.com/repos/ViRb3/wgcf/releases/latest \
| grep '"tag_name"' | cut -d'"' -f4)
193         [[ -z "$WGCF_VERSION" ]] && die "Could not fetch wgcf version. Check your internet connection."
194
195         sudo curl -sfl \
196             "https://github.com/ViRb3/wgcf/releases/download/${WGCF_VERSION}/wgcf_${WGCF_VERSION#v}
_linux_${WGCF_ARCH}" \
197             -o /usr/local/bin/wgcf

```

```

199     sudo chmod +x /usr/local/bin/wgcf
200 else
201     die "Unsupported OS. Install wgcf manually: https://github.com/ViRb3/wgcf/releases"
202 fi
203 fi
204
205 ok "wgcf is ready."
206
207 # 4. WARP profile
208 step "Generating Cloudflare WARP profile"
209
210 if [[ -f "wgcf-profile.conf" ]]; then
211     warn "wgcf-profile.conf already exists skipping key generation."
212 else
213     wgcf register --accept-tos 2>> output.log || die "wgcf register failed. Check your internet
214     connection."
215     wgcf generate 2>> output.log || die "wgcf generate failed."
216     ok "WARP profile generated."
217 fi
218
219 # Parse keys (IPv4 address only gluetun doesn't need the IPv6 one)
220 PRIVATE_KEY=$(awk '/PrivateKey/{print $3}' wgcf-profile.conf)
221 PUBLIC_KEY=$(awk '/PublicKey/{print $3}' wgcf-profile.conf)
222 ADDRESSES=$(awk '/Address/{print $3}' wgcf-profile.conf | grep -v ':' | head -1)
223
224 [[ -z "$PRIVATE_KEY" ]] && die "Could not parse PrivateKey from wgcf-profile.conf"
225 [[ -z "$PUBLIC_KEY" ]] && die "Could not parse PublicKey from wgcf-profile.conf"
226 [[ -z "$ADDRESSES" ]] && die "Could not parse IPv4 Address from wgcf-profile.conf"
227
228 ok "WARP keys extracted."
229
230 # 5. Tailscale auth key
231 step "Tailscale authentication"
232
233 TS_AUTHKEY=""
234 if [[ -f ".env" ]]; then
235     EXISTING_TS_KEY=$(read_env_var TS_AUTHKEY)
236     if [[ -n "$EXISTING_TS_KEY" ]]; then
237         TS_AUTHKEY="$EXISTING_TS_KEY"
238         warn "Found existing TS_AUTHKEY in .env. Skipping prompt."
239     fi
240 fi
241
242 if [[ -z "$TS_AUTHKEY" ]]; then
243     echo ""
244     echo "Generate a key at: https://login.tailscale.com/admin/settings/keys"
245     echo "'Generate auth key' enable Reusable if you plan to re-run setup"
246     read -rp "Paste your Tailscale auth key: " TS_AUTHKEY
247     echo ""
248 fi
249
250 [[ -z "$TS_AUTHKEY" ]] && die "Tailscale auth key is required."
251 [[ "$TS_AUTHKEY" != tskey-auth-* ]] && warn "Key doesn't look like a Tailscale auth key proceeding
252 anyway."
253
254 ok "Auth key accepted."
255
256 # 6. AdGuard Home admin password
257 step "AdGuard Home admin account"
258
259 # Hardcoded default credentials to prevent lockout
260 # Username: admin
261 # Password: nullexit
262 if [[ -f ".env" ]]; then
263     EXISTING_AG_PASS=$(read_env_var ADGUARD_PASSWORD)
264     if [[ -n "$EXISTING_AG_PASS" ]]; then
265         ADGUARD_PASSWORD="$EXISTING_AG_PASS"
266         ADGUARD_PWD_ESC="$EXISTING_AG_PASS"
267         ok "Found existing ADGUARD_PASSWORD in .env."
268     else
269         ADGUARD_PASSWORD="nullexit"
270         ADGUARD_PWD_ESC="nullexit"
271         ok "Default password set to 'nullexit'."
272     fi
273 else
274     ADGUARD_PASSWORD="nullexit"
275     ADGUARD_PWD_ESC="nullexit"
276     ok "Default password set to 'nullexit'."
277 fi

```



```

278 # 7. Write .env
279 step "Writing .env"
280
281 if [[ -f ".env" ]]; then
282     read -rp " .env already exists. Overwrite? [y/N]: " OVERWRITE
283     if [[ "$OVERWRITE" != "y" && "$OVERWRITE" != "Y" ]]; then
284         warn "Keeping existing .env."
285     else
286         WRITE_ENV=1
287     fi
288 else
289     WRITE_ENV=1
290 fi
291
292 if [[ "${WRITE_ENV:-0}" == "1" ]]; then
293     # Preserve existing configuration flags if .env already exists
294     EXISTING_PROFILE="medium"
295     EXISTING_MSS="1120"
296     EXISTING_HIJACK="true"
297     EXISTING_EXIT="true"
298     EXISTING_THRESH="6"
299     EXISTING_KILL="true"
300     EXISTING_BLOCKED="kp il"
301     EXISTING_HOST_MTU="1200"
302     EXISTING_SEED=""
303
304     if [[ -f ".env" ]]; then
305         [[ -n "$(read_env_var GATEWAY_RULE_PROFILE)" ]] && EXISTING_PROFILE=$(read_env_var
GATEWAY_RULE_PROFILE)
306         [[ -n "$(read_env_var GATEWAY_MSS)" ]] && EXISTING_MSS=$(read_env_var GATEWAY_MSS)
307         [[ -n "$(read_env_var GATEWAY_HIJACK_HOST)" ]] && EXISTING_HIJACK=$(read_env_var
GATEWAY_HIJACK_HOST)
308         [[ -n "$(read_env_var GATEWAY_USE_EXIT_NODE)" ]] && EXISTING_EXIT=$(read_env_var
GATEWAY_USE_EXIT_NODE)
309         [[ -n "$(read_env_var WARP_FAIL_THRESHOLD)" ]] && EXISTING_THRESH=$(read_env_var
WARP_FAIL_THRESHOLD)
310         [[ -n "$(read_env_var KILL_SWITCH)" ]] && EXISTING_KILL=$(read_env_var KILL_SWITCH)
311         [[ -n "$(read_env_var BLOCKED_COUNTRIES)" ]] && EXISTING_BLOCKED=$(read_env_var BLOCKED_COUNTRIES
)
312         [[ -n "$(read_env_var HOST_MTU)" ]] && EXISTING_HOST_MTU=$(read_env_var HOST_MTU)
313         [[ -n "$(read_env_var NULLEXIT_SEED)" ]] && EXISTING_SEED=$(read_env_var NULLEXIT_SEED)
314     fi
315
316     # README/devref both document NULLEXIT_SEED as "generated by setup.sh" actually
317     # generate it (rather than leaving the .env.example placeholder for the user to
318     # fill by hand), since crypto.sh --verify hard-fails toggle.sh's very first run
319     # without one, and there's nothing else in the setup flow that would create it.
320     if [[ -z "$EXISTING_SEED" ]]; then
321         EXISTING_SEED=$(openssl rand -hex 32)
322         ok "Generated NULLEXIT_SEED for script-integrity signing."
323     fi
324
325     cat > .env <<EOF
326 WIREGUARD_PRIVATE_KEY=${PRIVATE_KEY}
327 WIREGUARD_PUBLIC_KEY=${PUBLIC_KEY}
328 WIREGUARD_ADDRESSES=${ADDRESSES}
329 TS_AUTHKEY=${TS_AUTHKEY}
330 NULLEXIT_SEED=${EXISTING_SEED}
331 GATEWAY_RULE_PROFILE=${EXISTING_PROFILE}
332 # Safe MSS for double-tunneled traffic (Tailscale + WARP). Don't raise this
333 # devref.md 15.3.2 found 1180 (and even 1160) still too large and causes the
334 # exact mysterious mid-transfer stalls this value exists to prevent.
335 GATEWAY_MSS=${EXISTING_MSS}
336 # Lower the host Tailscale MTU to 1200 to prevent fragmentation when passing through the 1280-byte WARP
tunnel.
337 HOST_MTU=${EXISTING_HOST_MTU}
338 # Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own
internet is NOT hijacked).
339 GATEWAY_HIJACK_HOST=${EXISTING_HIJACK}
340 GATEWAY_USE_EXIT_NODE=${EXISTING_EXIT}
341 WARP_FAIL_THRESHOLD=${EXISTING_THRESH}
342 # PF fail-closed kill-switch. Leave true this is the core leak guarantee.
343 KILL_SWITCH=${EXISTING_KILL}
344 ADGUARD_PASSWORD=${ADGUARD_PASSWORD}
345 BLOCKED_COUNTRIES="${EXISTING_BLOCKED}"
346 EOF
347     ok ".env written."
348
349     # Sign the core scripts against the (possibly just-generated) seed now, so the
350     # very first toggle.sh run doesn't hard-fail crypto.sh --verify with no seed/no
351     # .signatures and no indication that a manual 'crypto.sh --sign' was needed.

```

```

352     if bash scripts/crypto.sh --sign >> output.log 2>&1; then
353         ok "Core scripts signed for integrity verification."
354     else
355         warn "Could not sign core scripts  run 'bash scripts/crypto.sh --sign' manually before toggle.sh."
356     fi
357 fi
358
359 # 8. Prepare directories
360 step "Preparing AdGuard directories"
361
362 mkdir -p adguard/conf adguard/work/userfilters
363
364 # Remove empty AdGuardHome.yaml that would cause a crash loop (same logic as Toggle-Gateway scripts)
365 if [[ -f "adguard/conf/AdGuardHome.yaml" ]] && ! -s "adguard/conf/AdGuardHome.yaml" ]]; then
366     rm "adguard/conf/AdGuardHome.yaml"
367     warn "Removed empty AdGuardHome.yaml to prevent crash loop."
368 fi
369
370 ok "Directories ready."
371
372 # 9. Compile DNS filter rules
373 step "Compiling DNS filter rules (this may take a minute)"
374
375 # 10. Start containers
376 step "Starting containers"
377 # Note: tailscale depends on adguardhome being healthy (port 5335 up),
378 # so Docker Compose will hold it back automatically until the wizard is done.
379 docker compose up -d
380 ok "Containers starting."
381
382 # 11. Wait for AdGuard Home setup wizard
383 step "Waiting for AdGuard Home setup wizard"
384 echo " (up to 60 seconds...)"
385
386 MAX=60; ELAPSED=0
387 until docker exec routing-fix curl -sf http://127.0.0.1:3000 >> output.log 2>&1; do
388     sleep 2; ELAPSED=$((ELAPSED + 2))
389     [[ $ELAPSED -ge $MAX ]] && die "AdGuard Home didn't come up.\nDebug with: docker compose logs adguardhome"
390 done
391
392 ok "AdGuard Home is up."
393
394 # 12. Configure AdGuard Home via API
395 step "Configuring AdGuard Home"
396
397 # Run all API calls in a single docker exec sh -c so the session cookie file
398 # is shared across calls without leaving the container.
399 docker exec routing-fix sh -c "
400     set -e
401
402     # Step 1: Complete the setup wizard (sets admin creds + DNS port to 5335)
403     curl -sf -X POST http://127.0.0.1:3000/control/install/configure \
404     -H 'Content-Type: application/json' \
405     -d '{\"web\":{\"ip\":\"0.0.0.0\",\"port\":3000},\"dns\":{\"ip\":\"0.0.0.0\",\"port\":5335},\"username\":\"admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
406     >/dev/null || true
407
408     # Step 2: Wait for AdGuard to restart after wizard
409     sleep 4
410     WAITED=0
411     until curl -sf http://127.0.0.1:3000 >/dev/null 2>&1; do
412         sleep 2; WAITED=$((WAITED + 2))
413         [ \ $WAITED -ge 30 ] && { echo 'AdGuard did not restart in time'; exit 1; }
414     done
415
416     # Step 3: Login to get session cookie
417     curl -sf -X POST http://127.0.0.1:3000/control/login \
418     -H 'Content-Type: application/json' \
419     -d '{\"name\":\"admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
420     -c /tmp/agh.cookie >/dev/null
421
422     # Step 4: Point upstream DNS back through the WARP tunnel
423     curl -sf -X POST http://127.0.0.1:3000/control/dns_config \
424     -H 'Content-Type: application/json' \
425     -b /tmp/agh.cookie \
426     -d '{\"upstream_dns\":[\"127.0.0.1:53\",\"[onion/]127.0.0.1:5353\"],\"bootstrap_dns\":[\"1.1.1.1\", \"8.8.8.8\"],\"upstream_mode\":\"load_balance\"}' \
427     >/dev/null
428

```

```

429
430 # Step 5: Register the compiled rules file as a filter list
431 curl -sf -X POST http://127.0.0.1:3000/control/filtering/add_url \
432     -H 'Content-Type: application/json' \
433     -b /tmp/agh.cookie \
434     -d '{"name":"'Custom Compiled Rules',"url":"'"/opt/adguardhome/work/userfilters/compiled_rules.
435         txt\":"' \
436     >/dev/null || true
437
438 # Step 6: Force a filter refresh to load the rules immediately
439 curl -sf -X POST http://127.0.0.1:3000/control/filtering/refresh \
440     -H 'Content-Type: application/json' \
441     -b /tmp/agh.cookie \
442     -d '{"whitelist":false}' \
443     >/dev/null || true
444
445 rm -f /tmp/agh.cookie
446
447 "
448 ok "AdGuard Home configured."
449
450 # 13. Wait for Tailscale
451 step "Waiting for Tailscale to authenticate"
452 echo " (Tailscale only starts once AdGuard's healthcheck passes up to 90s)"
453
454 MAX=90; ELAPSED=0; TS_IP=""
455 until [[ -n "$TS_IP" ]]; do
456     TS_IP=$(docker exec tailscale tailscale ip -4 2>> output.log || true)
457     if [[ -z "$TS_IP" ]]; then
458         sleep 3; ELAPSED=$((ELAPSED + 3))
459         if [[ $ELAPSED -ge $MAX ]]; then
460             warn "Tailscale hasn't authenticated yet."
461             warn "Once it does, re-run the toggle script it will resolve the IP dynamically."
462             break
463         fi
464     fi
465 done
466
467 if [[ -n "$TS_IP" ]]; then
468     ok "Tailscale IP: ${TS_IP} (resolved dynamically by the toggle script on each startup)"
469 fi

```

53 scripts/sweep.sh

```

1 #!/bin/bash
2 # scripts/sweep.sh nullexit gateway health sweep
3 #
4 # Automates the 7-check post-restart health sweep (sweep.md 1 + tailnet ping
5 # + throughput). It answers one question: "after a toggle/restart/wake, is
6 # the gateway actually healthy end-to-end no leak, direct P2P, DNS
7 # filtering live, kill-switch armed, and fast enough to actually use?" It is
8 # READ-ONLY with respect to gateway state: it never mutates routing, PF, DNS,
9 # or containers (the throughput check only pulls a test file over curl the
10 # same kind of traffic any app already sends, nothing is exec'd into or
11 # installed in the container), so it is safe to run against a live gateway at
12 # any time (including from a remote session see the caveat in agent.md L16).
13 #
14 # The 7 checks (1-5 mirror sweep.md 1):
15 # 1. Containers warp, adguardhome, tailscale, tor, routing-fix up/healthy
16 # 2. Double-tunnel/leak container WARP exit IP == host exit IP, both warp=on
17 # 3. Hostcontainer path Tailscale exit-node is a DIRECT P2P conn, not DERP relay
18 # 4. AdGuard filtering compiled rule counts present + live DNS interception
19 # 5. PF kill-switch pfctl Enabled + anchor com.apple/nullexit carries rules
20 # 6. Tailnet reachability every ONLINE Tailscale peer answers a tailscale ping
21 # (DERP-relayed peers are flagged as throughput-degraded, not a failure)
22 # 7. Throughput host path vs. WARP-only baseline (via the container's
23 # SOCKS5 proxy, no exec/install inside the container). Catches "slow AND
24 # dropping mid-transfer" that a one-shot ping/latency check can't see.
25 # Real per-device (phone) throughput can't be measured without an agent
26 # running on the device, so mesh peers get a path/latency flag instead
27 # (see check 6).
28 #
29 # A timestamped report is written to the project root (one file per run, so
30 # runs can be diffed). stdout carries a coloured summary + a final PASS/FAIL
31 # verdict; exit code is 0 when every check passes, 1 otherwise (so it can gate
32 # CI / launchd / a post-restart hook).

```

```

33 #
34 # Usage:
35 #   bash scripts/sweep.sh           # run the sweep, write report, print verdict
36 #   bash scripts/sweep.sh --quiet   # only the final verdict line + non-zero on FAIL
37 #   bash scripts/sweep.sh --help    # this message
38
39 set -uo pipefail
40 # NOTE: errexit is intentionally OFF. Several probes are EXPECTED to fail on an
41 # unhealthy gateway and that failure IS the diagnostic signal we record each
42 # exit status rather than letting one abort the whole sweep.
43
44 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
45 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
46 TIMESTAMP="$(date -u +%Y%m%dT%H%M%S)"
47 cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
48 source "$SCRIPT_DIR/common.sh"
49 LOG_FILE="$PROJECT_ROOT/output.log"
50 REPORT_FILE="$PROJECT_ROOT/sweep-$TIMESTAMP.txt"
51
52 # Argument parsing
53 QUIET=false
54 case "${1:-}" in
55   --quiet|-q) QUIET=true ;;
56   --help|-h)  sed -n '2,27p' "$0"; exit 0 ;;
57   *)          : ;;
58   *)          echo "Unknown argument: $1" >&2
59               echo "Run with --help for usage." >&2
60               exit 2 ;;
61 esac
62
63 # Colours (auto-disabled when stdout is not a tty, e.g. piped)
64 if [ -t 1 ]; then
65   RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
66   BOLD='\033[1m'; NC='\033[0m'
67 else
68   RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
69 fi
70
71 # Homebrew on PATH (same as toggle.sh / recover.sh / diagnose-host-leak)
72 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
73
74 # Output helpers write to BOTH stdout and the report file
75 exec 3>> "$REPORT_FILE" || {
76   printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
77   exit 1
78 }
79 emit() { [ "$QUIET" = true ] || printf '%b\n' "$1"; printf '%b\n' "$1" | sed 's/\x1b\[[0-9;]*m//g' >&3; }
80 section() { emit ""; emit "${BOLD} $1 ${NC}"; }
81 line() { emit " $1"; }
82 ok() { line "${GREEN}${NC} $1"; }
83 warn() { line "${YELLOW}${NC} $1"; }
84 fail() { line "${RED}${NC} $1"; }
85
86 # common.sh redefines run_with_timeout; it is a pure-bash macOS-safe timeout
87 # (GNU 'timeout' is absent on stock macOS). We reuse the sourced version.
88
89 # Result accounting
90 # Each check appends "PASS"/"WARN"/"FAIL" so the verdict is a pure roll-up.
91 declare -a RESULTS=()
92 record() { RESULTS+=("${1}|${2}"); } # status|label
93
94 # Header
95 if [ "$QUIET" != true ]; then
96   emit ""
97   emit ""
98   emit " n u l l e x i t   h e a l t h   s w e e p "
99   emit ""
100   emit " Timestamp (UTC): $TIMESTAMP"
101   emit " Project root:    $PROJECT_ROOT"
102   emit " Report file:     $REPORT_FILE"
103   emit ""
104 fi
105
106 GATEWAY_TS_IP="$(read_adguard_ip)" # gateway Tailscale IP from .gateway_ip
107
108 #
109 section "1/7 Containers"
110 #
111 # All five services must be running; warp + adguardhome additionally expose a
112 # healthcheck and must be 'healthy'. 'docker compose ps' is the source of truth.

```

```

113 EXPECTED_SERVICES="warp adguardhome tailscale tor routing-fix"
114 if ! command -v docker >/dev/null 2>&1; then
115     fail "docker CLI not on PATH cannot inspect containers"
116     record FAIL "containers"
117 else
118     PS_OUT=$(run_with_timeout 20 docker compose ps 2>>"$LOG_FILE")
119     missing=""; unhealthy=""
120     for svc in $EXPECTED_SERVICES; do
121         row=$(printf '%s\n' "$PS_OUT" | awk -v s="$svc" ' $1==s || $0 ~ ("[:space:]")s"[:space:]')
122         if [ -z "$row" ] || ! printf '%s' "$row" | grep -qiE 'up|running'; then
123             missing="$missing $svc"
124             elif printf '%s' "$row" | grep -qi 'unhealthy'; then
125                 unhealthy="$unhealthy $svc"
126             fi
127         done
128     if [ -z "$missing" ] && [ -z "$unhealthy" ]; then
129         ok "all 5 services up (warp/adguardhome healthy, tailscale/tor/routing-fix running)"
130         record PASS "containers"
131     else
132         [ -n "$missing" ] && fail "not up/running:$missing"
133         [ -n "$unhealthy" ] && fail "reporting unhealthy:$unhealthy"
134         record FAIL "containers"
135     fi
136     line " docker compose ps "
137     printf '%s\n' "$PS_OUT" | awk '{print " " "$0}" >&3
138     [ "$QUIET" = true ] || printf '%s\n' "$PS_OUT" | awk '{print " " "$0}'
139 fi
140
141 #
142 section "2/7 Double-tunnel / IP-leak"
143 #
144 # The container's WARP egress and the host's egress must be the SAME public IP
145 # with warp=on for both. Equal IPs all host traffic is exiting through the
146 # same WARP tunnel the container uses (no split / raw-ISP leak).
147 # The warp container (gluetun) ships wget, not curl fetch the trace from
148 # inside its netns with whichever client is present.
149 CTR_TRACE=$(run_with_timeout 12 docker compose exec -T warp sh -c \
150     'curl -s --max-time 8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
151     || wget -qO- --timeout=8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null' \
152     2>>"$LOG_FILE")
153 HOST_TRACE=$(run_with_timeout 12 curl -s --max-time 8 \
154     https://www.cloudflare.com/cdn-cgi/trace 2>>"$LOG_FILE")
155 CTR_IP=$(printf '%s' "$CTR_TRACE" | awk -F= '/^ip=/ {print $2; exit}')
156 CTR_WARP=$(printf '%s' "$CTR_TRACE" | awk -F= '/^warp=/ {print $2; exit}')
157 HOST_IP=$(printf '%s' "$HOST_TRACE" | awk -F= '/^ip=/ {print $2; exit}')
158 HOST_WARP=$(printf '%s' "$HOST_TRACE" | awk -F= '/^warp=/ {print $2; exit}')
159 line "container exit: ip=${CTR_IP:-?} warp=${CTR_WARP:-?}"
160 line "host exit: ip=${HOST_IP:-?} warp=${HOST_WARP:-?}"
161 if [ -z "$CTR_IP" ] || [ -z "$HOST_IP" ]; then
162     fail "could not obtain both egress IPs (container or host trace failed)"
163     record FAIL "leak"
164 elif [ "$CTR_IP" = "$HOST_IP" ] && [ "$HOST_WARP" = "on" ]; then
165     ok "no leak host egress == container egress ($HOST_IP), warp=on"
166     record PASS "leak"
167 elif [ "$HOST_WARP" != "on" ]; then
168     fail "host warp=$HOST_WARP host traffic is NOT going through WARP (LEAK)"
169     record FAIL "leak"
170 else
171     fail "egress mismatch: host $HOST_IP != container $CTR_IP split/leak"
172     record FAIL "leak"
173 fi
174
175 #
176 section "3/7 Host container path (direct P2P, not DERP relay)"
177 #
178 # 'tailscale status' annotates each peer with 'direct <ip:port>' or
179 # 'relay <derp>'. The gateway peer should be DIRECT a DERP relay hop adds
180 # latency and signals NAT traversal failure between host and gateway.
181 if ! command -v tailscale >/dev/null 2>&1; then
182     fail "tailscale CLI not on PATH"
183     record FAIL "p2p"
184 else
185     TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
186     PEER_LINE=$(printf '%s\n' "$TS_OUT" | grep -E "(^|[:space:])$GATEWAY_TS_IP" | head -1)
187     if [ -z "$PEER_LINE" ]; then
188         warn "gateway $GATEWAY_TS_IP not found in host peer list"
189         line "$(printf '%s\n' "$TS_OUT" | head -5 | awk '{print " " "$0}')"
190         record WARN "p2p"
191     elif printf '%s' "$PEER_LINE" | grep -q 'direct '; then
192         conn=$(printf '%s' "$PEER_LINE" | grep -oE 'direct [0-9.]+:[0-9]+')
193         ok "direct P2P to gateway $conn (no DERP relay)"

```

```

194     record PASS "p2p"
195 elif printf '%s' "$PEER_LINE" | grep -q 'relay'; then
196     relay=$(printf '%s' "$PEER_LINE" | grep -oE 'relay "[^"]+"' )
197     warn "gateway reached via DERP $relay (not direct) NAT traversal degraded"
198     record WARN "p2p"
199 else
200     warn "gateway peer state indeterminate (idle/no active conn)"
201     line "      $PEER_LINE"
202     record WARN "p2p"
203 fi
204 fi
205
206 #
207 section "4/7  AdGuard filtering"
208 #
209 # Two signals: (a) the compiled ruleset exists with a non-trivial block count,
210 # (b) the gateway resolver actually answers a query (live interception).
211 COMPILED="adguard/work/userfilters/compiled_rules.txt"
212 if [ -f "$COMPILED" ]; then
213     TOTAL=$(grep -m1 '! Total Block Rules:' "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
214     NATIVE=$(grep -m1 'Native AdGuard Rules:' "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
215     if [ -n "$TOTAL" ] && [ "$TOTAL" -gt 0 ] 2>/dev/null; then
216         ok "compiled ruleset present Total Block Rules: $TOTAL / Native AdGuard: ${NATIVE:-?}"
217         RULES_OK=true
218     else
219         warn "compiled_rules.txt present but block-rule count unreadable"
220         RULES_OK=false
221     fi
222 else
223     warn "compiled_rules.txt missing (rule-compile may not have run)"
224     RULES_OK=false
225 fi
226 # Live interception: the gateway resolver must answer (dig via .gateway_ip).
227 DNS_OK=false
228 if command -v dig >/dev/null 2>&1 && [ -n "$GATEWAY_TS_IP" ]; then
229     DIG_OUT=$(run_with_timeout 8 dig +time=3 +tries=1 @$GATEWAY_TS_IP google.com A 2>>"$LOG_FILE")
230     if printf '%s' "$DIG_OUT" | grep -qE 'ANSWER SECTION|status: NOERROR'; then
231         ok "live DNS interception gateway $GATEWAY_TS_IP resolved google.com"
232         DNS_OK=true
233     else
234         warn "gateway $GATEWAY_TS_IP did not answer a test query"
235     fi
236 else
237     warn "skipped live-resolution probe (dig missing or gateway IP unknown)"
238 fi
239 if [ "$RULES_OK" = true ] && [ "$DNS_OK" = true ]; then
240     record PASS "adguard"
241 elif [ "$RULES_OK" = true ] || [ "$DNS_OK" = true ]; then
242     record WARN "adguard"
243 else
244     record FAIL "adguard"
245 fi
246
247 #
248 section "5/7  PF kill-switch"
249 #
250 # PF must be Enabled AND the com.apple/nullexit anchor must carry rules.
251 # An empty/absent anchor means the fail-closed lane is not actually armed.
252 PF_STATUS=$(run_with_timeout 6 sudo -n pfctl -s info 2>>"$LOG_FILE" | awk '/^Status:/{print $2;exit}')
253 ANCHOR_RULES=$(run_with_timeout 6 sudo -n pfctl -a com.apple/nullexit -sr 2>>"$LOG_FILE")
254 if [ "$PF_STATUS" = "Enabled" ] && [ -n "$ANCHOR_RULES" ]; then
255     RULE_N=$(printf '%s\n' "$ANCHOR_RULES" | grep -cvE '\s*')
256     ok "PF Enabled + anchor com.apple/nullexit armed ($RULE_N rules)"
257     line "$(printf '%s\n' "$ANCHOR_RULES" | head -4 | awk '{print "      "$0}')"
258     record PASS "pf"
259 elif [ "$PF_STATUS" = "Enabled" ]; then
260     fail "PF Enabled but anchor com.apple/nullexit has NO rules kill-switch not armed"
261     record FAIL "pf"
262 elif [ -z "$PF_STATUS" ]; then
263     warn "could not read PF status (needs passwordless sudo for 'pfctl -s info')"
264     record WARN "pf"
265 else
266     fail "PF Status: $PF_STATUS kill-switch disabled"
267     record FAIL "pf"
268 fi
269
270 #
271 section "6/7  Tailnet reachability (ping all online devices)"
272 #
273 # Every peer the coordination server reports as ONLINE should answer a
274 # 'tailscale ping' (TSMP probes the WireGuard data path; ICMP is avoided since

```

```

275 # phones/laptops often drop it). Offline peers and this host are skipped.
276 # An online-but-unreachable peer means control plane and data path disagree.
277 if ! command -v tailscale >/dev/null 2>&1; then
278     fail "tailscale CLI not on PATH"
279     record FAIL "tailnet"
280 else
281     SELF_TS_IP=$(run_with_timeout 6 tailscale ip -4 2>>"$LOG_FILE")
282     [ -n "${TS_OUT:-}" ] || TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
283     # Online peers = rows starting with a Tailscale IP, minus self and 'offline'.
284     # (A '-' status column means online-but-idle, so it is kept.)
285     PEERS=$(printf '%s\n' "$TS_OUT" | awk -v self="$SELF_TS_IP" \
286         '$1 ~ /^100\./ && $1 != self && $0 !~ /^([:space:]]offline(,|[:space:]))$/ {print $1, $2
287     }')
288     if [ -z "$PEERS" ]; then
289         warn "no online peers to ping (everything else in the tailnet is offline)"
290         record WARN "tailnet"
291     else
292         unreached=0
293         while read -r peer_ip peer_name; do
294             PONG=$(run_with_timeout 10 tailscale ping --timeout=3s --c 1 --until-direct=false "$peer_ip" 2>>"$LOG_FILE")
295             if printf '%s' "$PONG" | grep -q '^pong'; then
296                 detail=$(printf '%s\n' "$PONG" | head -1 | sed -E 's/^pong from [^ ]+ \([^\)]+\) //' )
297                 if printf '%s' "$detail" | grep -q 'DERP'; then
298                     ok "$peer_name ($peer_ip) pong $detail relayed, expect degraded throughput (see check 7)"
299                 else
300                     ok "$peer_name ($peer_ip) pong $detail"
301                 fi
302             else
303                 fail "$peer_name ($peer_ip) no pong (online per control plane, data path dead)"
304                 unreached=$((unreached+1))
305             fi
306         done <<< "$PEERS"
307         n_peers=$(printf '%s\n' "$PEERS" | wc -l | tr -d '[:space:]')
308         if [ "$unreached" -eq 0 ]; then
309             record PASS "tailnet"
310         elif [ "$unreached" -lt "$n_peers" ]; then
311             record WARN "tailnet"
312         else
313             record FAIL "tailnet"
314         fi
315     fi
316 fi
317
318 #
319 # section "7/7 Throughput (host path vs. WARP-only baseline)"
320 #
321 # A single ping only proves the path is UP, not that it stays up under real
322 # load devref.md 15.3.1 documents sustained-transfer bufferbloat/reset
323 # behavior a quick latency check can't see. Pulls a real test file twice:
324 # once over the host's normal route (full double-tunnel: Tailscale exit-node
325 # wrap + WARP) and once straight through the warp container's own SOCKS5
326 # proxy (WARP only, no Tailscale wrap) port read from the live
327 # /tmp/nullexit-ports.env, so this never execs into or installs anything in
328 # the container (sweep stays read-only w.r.t. gateway state). Comparing the
329 # two shows whether slowness/resets come from WARP itself or from the extra
330 # Tailscale-wrap hop. Skip with SWEEP_SKIP_THROUGHPUT=true in .env.
331 #
332 # Target: a fast, high-bandwidth CDN, not a slow/distant one. speedtest.net-style
333 # test hosts (e.g. speedtest.tele2.net) were tried first and turned out to be the
334 # bottleneck themselves even RAW, gateway-off ISP speed to tele2.net was ~3.2
335 # Mbps, vs ~32 Mbps raw to a fast CDN so a slow test target would have silently
336 # blamed nullexit for a benchmark artifact. Cloudflare's own speed-test endpoint
337 # 403s WARP egress IPs specifically (abuse-prevention on shared/CGNAT IPs), so
338 # Hetzner (also fast, does not flag WARP IPs) is used instead.
339 THROUGHPUT_SKIP=$(read_env_var "SWEEP_SKIP_THROUGHPUT" | tr '[:upper:]' '[:lower:]')
340 if [ "$THROUGHPUT_SKIP" = "true" ]; then
341     warn "skipped (SWEEP_SKIP_THROUGHPUT=true)"
342     record WARN "throughput"
343 else
344     TP_URL="${SWEEP_THROUGHPUT_URL:-https://ash-speed.hetzner.com/100MB.bin}"
345     TP_TIMEOUT="${SWEEP_THROUGHPUT_TIMEOUT:-30}"
346
347     measure() {
348         # $1 = extra curl args (word-split; may be empty)
349         # curl's own --max-time must win the race against run_with_timeout's outer
350         # SIGTERM watcher a process killed by signal never reaches its own -w
351         # output, so if both timers were equal, the watcher could kill curl right
352         # before it prints stats, and every leg would misreport as "no response".
353         # The outer timeout is a dead-man's-switch (+5s), not the real budget.
354         run_with_timeout $((TP_TIMEOUT + 5)) curl -sS --max-time "$TP_TIMEOUT" $1 -o /dev/null \
355             -w 'code=%{http_code} size=%{size_download} time=%{time_total} bps=%{speed_download}' \

```



```

354     "$TP_URL" 2>>"$LOG_FILE"
355 }
356 leg_verdict() { # $1=curl_rc $2=http_code -> ok|stall|blocked|unknown
357 [ -z "$2" ] && { echo unknown; return; }
358 [ "$2" != "200" ] && { echo blocked; return; }
359 case "$1" in
360 0|28) echo ok ;; # 0=completed; 28=hit our own --max-time while still healthy
361 56) echo stall ;; # curl's actual "Recv failure: Connection reset by peer"
362 *) echo unknown ;;
363 esac
364 }
365 to_mbps() { awk -v b="$1:-" 'BEGIN{ if (b=="") print "?"; else printf "%.1f", (b*8)/1000000 }'; }
366
367 HOST_TP=$(measure ""); HOST_RC=$?
368 HOST_CODE=$(printf '%s' "$HOST_TP" | grep -oE 'code=[0-9]+' | cut -d= -f2)
369 HOST_SIZE=$(printf '%s' "$HOST_TP" | grep -oE 'size=[0-9]+' | cut -d= -f2)
370 HOST_TIME=$(printf '%s' "$HOST_TP" | grep -oE 'time=[0-9.]+' | cut -d= -f2)
371 HOST_BPS=$(printf '%s' "$HOST_TP" | grep -oE 'bps=[0-9.]+' | cut -d= -f2)
372 HOST_VERDICT=$(leg_verdict "$HOST_RC" "$HOST_CODE")
373 HOST_MBPS=$(to_mbps "$HOST_BPS")
374 case "$HOST_VERDICT" in
375 ok) ok "host path: ${HOST_MBPS} Mbps (${HOST_SIZE:-0} bytes in ${HOST_TIME:-?}s)" ;;
376 stall) warn "host path: connection reset mid-transfer after ${HOST_SIZE:-0} bytes sustained-load
drop, see devref.md 15.3.1" ;;
377 blocked) warn "host path: test endpoint returned HTTP ${HOST_CODE:-?} (egress IP likely flagged)
inconclusive, not a gateway fault" ;;
378 unknown) warn "host path: no response (timeout after ${TP_TIMEOUT}s)" ;;
379 esac
380
381 SOCKS_PROXY_PORT=""
382 [ -f /tmp/nullexit-ports.env ] && SOCKS_PROXY_PORT=$(grep -oE 'SOCKS_PROXY_PORT=[0-9]+' /tmp/nullexit-
ports.env | cut -d= -f2)
383 if [ -z "$SOCKS_PROXY_PORT" ]; then
384 warn "WARP-only baseline: SOCKS_PROXY_PORT unknown (no /tmp/nullexit-ports.env gateway not started
via toggle.sh?)"
385 CTR_VERDICT=unknown
386 else
387 CTR_TP=$(measure "--socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT"); CTR_RC=$?
388 CTR_CODE=$(printf '%s' "$CTR_TP" | grep -oE 'code=[0-9]+' | cut -d= -f2)
389 CTR_SIZE=$(printf '%s' "$CTR_TP" | grep -oE 'size=[0-9]+' | cut -d= -f2)
390 CTR_TIME=$(printf '%s' "$CTR_TP" | grep -oE 'time=[0-9.]+' | cut -d= -f2)
391 CTR_BPS=$(printf '%s' "$CTR_TP" | grep -oE 'bps=[0-9.]+' | cut -d= -f2)
392 CTR_VERDICT=$(leg_verdict "$CTR_RC" "$CTR_CODE")
393 CTR_MBPS=$(to_mbps "$CTR_BPS")
394 case "$CTR_VERDICT" in
395 ok) ok "WARP-only baseline: ${CTR_MBPS} Mbps (${CTR_SIZE:-0} bytes in ${CTR_TIME:-?}s)" ;;
396 stall) warn "WARP-only baseline: connection reset mid-transfer after ${CTR_SIZE:-0} bytes bug is
in WARP itself, not the Tailscale wrap" ;;
397 blocked) warn "WARP-only baseline: test endpoint returned HTTP ${CTR_CODE:-?} inconclusive" ;;
398 unknown) warn "WARP-only baseline: no response (timeout after ${TP_TIMEOUT}s)" ;;
399 esac
400 fi
401
402 if [ "$HOST_VERDICT" = "ok" ] && [ "$CTR_VERDICT" = "ok" ]; then
403 if [ "$HOST_MBPS" != "?" ] && [ "$CTR_MBPS" != "?" ]; then
404 DELTA=$(awk -v h="$HOST_MBPS" -v c="$CTR_MBPS" 'BEGIN{ if (c<=0) print 0; else printf "%.0f", 100*(
c-h)/c }')
405 [ "$DELTA" -ge 30 ] 2>/dev/null && warn "host is ${DELTA}% slower than the WARP-only baseline the
Tailscale exit-node wrap is the added cost here, not WARP"
406 fi
407 record PASS "throughput"
408 elif [ "$HOST_VERDICT" = "stall" ] || [ "$CTR_VERDICT" = "stall" ]; then
409 record WARN "throughput"
410 elif [ "$HOST_VERDICT" = "unknown" ] && [ "$CTR_VERDICT" = "unknown" ]; then
411 record FAIL "throughput"
412 else
413 record WARN "throughput"
414 fi
415 fi
416
417 #
418 # Cover traffic (opt-in) informational only, NOT a scored check.
419 # Shown solely when NOISE_ENABLED=true so the sweep stays read-only and the
420 # 6-check verdict is unchanged. Reports whether the padding daemon is alive.
421 NOISE_ON=$(read_env_var "NOISE_ENABLED" | tr '[:upper:]' '[:lower:]')
422 if [ "$NOISE_ON" = "true" ]; then
423 section "Cover traffic (opt-in, informational)"
424 NOISE_PID_FILE="/tmp/nullexit-noise.pid"
425 npid=$(cat "$NOISE_PID_FILE" 2>/dev/null)
426 if [ -n "$npid" ] && kill -0 "$npid" 2>/dev/null; then
427 ok "dummy-padding daemon running (pid $npid) see 'bash scripts/noise.sh status'"

```

```

428     else
429         warn "NOISE_ENABLED=true but no padding daemon running start it: bash scripts/noise.sh start"
430     fi
431 fi
432
433 #
434 # DNS anomaly scan (informational, read-only, NOT a scored check the 6-check
435 # verdict is unchanged). It only READS the AdGuard querylog; never blocks DNS.
436 # It FAILS LOUD: a broken detector (crash, corrupt model, empty log, failing
437 # self-test) shows a red with the reason it is NEVER silently reported as
438 # clean, and an untrained detector is announced rather than skipped in silence.
439 QUERYLOG="$PROJECT_ROOT/adguard/work/data/querylog.json"
440 DNS_TOOL="$SCRIPT_DIR/dns_anomaly_detector.py"
441 if command -v python3 >/dev/null 2>&1 && [ -f "$DNS_TOOL" ]; then
442     section "DNS anomaly scan (informational)"
443     # 1) Prove the detection logic itself works before trusting any "clean".
444     ST_ERR=$(python3 "$DNS_TOOL" selftest 2>&1 >/dev/null); ST_RC=$?
445     if [ "$ST_RC" -ne 0 ]; then
446         fail "DNS detector SELF-TEST FAILED (rc=$ST_RC) detector is broken, do NOT trust results: ${ST_ERR:-
447             see stderr}"
448     elif [ ! -f "$PROJECT_ROOT/.dns_baseline.json" ]; then
449         warn "DNS detector NOT trained (no .dns_baseline.json) not protecting; run: python3 scripts/
450             dns_anomaly_detector.py learn"
451     elif [ ! -f "$QUERYLOG" ]; then
452         warn "DNS detector trained but querylog absent ($QUERYLOG) nothing to scan"
453     else
454         # 2) Real scan. Capture stderr + exit code; do NOT swallow failures.
455         DNS_ERR=$(python3 "$DNS_TOOL" scan --quiet 2>&1 >/tmp/nullexit-dns-scan.out); DNS_RC=$?
456         DNS_OUT=$(cat /tmp/nullexit-dns-scan.out 2>/dev/null); rm -f /tmp/nullexit-dns-scan.out
457         case "$DNS_RC" in
458             0) ok "no DNS-tunneling / DGA signatures $DNS_OUT" ;;
459             1) warn "$DNS_OUT review: python3 scripts/dns_anomaly_detector.py scan" ;;
460             *) fail "DNS scan FAILED to run (rc=$DNS_RC) NOT a clean result: ${DNS_ERR:-unknown error}" ;;
461         esac
462     fi
463 fi
464
465 #
466 section "VERDICT"
467 #
468 pass=0; warnc=0; failc=0
469 for r in "${RESULTS[@]"; do
470     case "${r%[*]}" in PASS) pass=$((pass+1));; WARN) warnc=$((warnc+1));; FAIL) failc=$((failc+1));; esac
471 done
472 for r in "${RESULTS[@]"; do
473     st="${r%[*]"; lbl="${r##[*]}"
474     case "$st" in
475         PASS) ok "$lbl" ;;
476         WARN) warn "$lbl" ;;
477         FAIL) fail "$lbl" ;;
478     esac
479 done
480 emit ""
481 if [ "$failc" -eq 0 ] && [ "$warnc" -eq 0 ]; then
482     emit "${GREEN}${BOLD}SWEEP PASS gateway healthy (${pass}/${#RESULTS[@]} checks green){NC}"
483     EXIT=0
484 elif [ "$failc" -eq 0 ]; then
485     emit "${YELLOW}${BOLD}SWEEP PASS with warnings ${pass} pass, ${warnc} warn, 0 fail${NC}"
486     EXIT=0
487 else
488     emit "${RED}${BOLD}SWEEP FAIL ${pass} pass, ${warnc} warn, ${failc} fail${NC}"
489     EXIT=1
490 fi
491 emit ""
492 [ "$QUIET" = true ] || emit "Full report: $REPORT_FILE"
493 exit "$EXIT"

```

54 scripts/unlock-files.sh

```

1 #!/bin/bash
2 # Resolve project root relative to this script's location, so the script
3 # works regardless of the caller's CWD.
4 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
5 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
6
7 echo "Unlocking .env..."

```

```

8 if [ -f "$PROJECT_ROOT/.env" ]; then
9     sudo cat "$PROJECT_ROOT/.env" > "$PROJECT_ROOT/.env.tmp" && mv "$PROJECT_ROOT/.env.tmp" "$PROJECT_ROOT
    /.env"
10     echo ".env unlocked"
11 fi
12
13 echo "Unlocking AdGuard cache files..."
14 for f in "$PROJECT_ROOT"/adguard/work/userfilters/compiled_rules.txt \
15         "$PROJECT_ROOT"/adguard/work/userfilters/ip_blocklist.ipset \
16         "$PROJECT_ROOT"/adguard/work/userfilters/cache/*.txt \
17         "$PROJECT_ROOT"/adguard/work/userfilters/cache/ip/*.txt \
18         "$PROJECT_ROOT"/adguard/work/data/filters/*.txt; do
19     if [ -f "$f" ]; then
20         cat "$f" > "$f.tmp" && mv "$f.tmp" "$f"
21         echo "Unlocked $f"
22     fi
23 done
24
25 echo ""
26 echo "All locked files have been successfully bypassed via atomic rename!"

```

55 tor-bridges.txt

56 white_list.txt

```

1 0.client-channel.google.com
2 1drv.com
3 2.android.pool.ntp.org
4 akamaihd.net
5 akamaitechnologies.com
6 akamaized.net
7 amazonaws.com
8 android.clients.google.com
9 api-adservices.apple.com
10 api.ipify.org
11 api.rlje.net
12 app-api.ted.com
13 appleid.apple.com
14 apps.skype.com
15 appsbackup-pa.clients6.google.com
16 appsbackup-pa.googleapis.com
17 appspot-preview.l.google.com
18 apt.sonarr.tv
19 aspnetcdn.com
20 attestation.xboxlive.com
21 ax.phobos.apple.com.edgesuite.net
22 books-analytics-events.apple.com
23 brightcove.net
24 c.s-microsoft.com
25 cdn.cloudflare.net
26 cdn.embedly.com
27 cdn.optimizely.com
28 cdn.vidible.tv
29 cdn2.optimizely.com
30 cdn3.optimizely.com
31 cdnjs.cloudflare.com
32 cert.mgt.xboxlive.com
33 clientconfig.passport.net
34 clients1.google.com
35 clients2.google.com
36 clients3.google.com
37 clients4.google.com
38 clients5.google.com
39 clients6.google.com
40 connectivitycheck.android.com
41 connectivitycheck.gstatic.com
42 continuum.dds.microsoft.com
43 cpms.spop10.ams.plex.bz
44 cpms35.spop10.ams.plex.bz
45 cse.google.com
46 ctldl.windowsupdate.com

```

47 cws.conviva.com
48 d2c8v521l5s99u.cloudfront.net
49 d2gatte9o95jao.cloudfront.net
50 dashboard.plex.tv
51 dataplicity.com
52 def-vef.xboxlive.com
53 delivery.vidible.tv
54 dev.virtualearth.net
55 device.auth.xboxlive.com
56 display.ugc.bazaarvoice.com
57 displaycatalog.mp.microsoft.com
58 dl.delivery.mp.microsoft.com
59 dl.dropbox.com
60 dl.dropboxusercontent.com
61 dns.msftncsi.com
62 download.sonarr.tv
63 drift.com
64 driftt.com
65 dynupdate.no-ip.com
66 ecn.dev.virtualearth.net
67 edge.api.brightcove.com
68 eds.xboxlive.com
69 fonts.gstatic.com
70 forums.sonarr.tv
71 g.live.com
72 geo-prod.do.dsp.mp.microsoft.com
73 geo3.ggpht.com
74 gfwsl.geforce.com
75 giphy.com
76 github.com
77 github.io
78 googleapis.com
79 gravatar.com
80.gstatic.com
81 help.ui.xboxlive.com
82 hls.ted.com
83 i.ytimg.com
84 il.ytimg.com
85 iadsdk.apple.com
86 imagesak.secureserver.net
87 img.vidible.tv
88 imgix.net
89 imgs.xkcd.com
90 instantmessaging-pa.googleapis.com
91 intercom.io
92 jquery.com
93 jsdelivr.net
94 keystone.mwbsys.com
95 lastfm-img2.akamaized.net
96 licensing.xboxlive.com
97 live.com
98 livepassdl.conviva.com
99 login.live.com
100 login.microsoftonline.com
101 manifest.googlevideo.com
102 meta-db-worker02.pop.ric.plex.bz
103 meta.plex.bz
104 meta.plex.tv
105 microsoftonline.com
106 msftncsi.com
107 my.plexapp.com
108 nexusrules.officeapps.live.com
109 nine.plugins.plexapp.com
110 no-ip.com
111 node.plexapp.com
112 notes-analytics-events.apple.com
113 notify.xboxlive.com
114 npr-news.streaming.adswizz.com
115 ns1.dropbox.com
116 ns2.dropbox.com
117 ol.email.plex.tv
118 o2.sg0.plex.tv
119 ocsp.apple.com
120 office.com
121 office.net
122 office365.com
123 officeclient.microsoft.com
124 om.cbsi.com
125 onedrive.live.com
126 outlook.live.com
127 outlook.office365.com

```

128 pings.conviva.com
129 placeholder.it
130 placeholder.imgix.net
131 players.brightcove.net
132 pricelist.skype.com
133 products.office.com
134 proxy.plex.bz
135 proxy.plex.tv
136 proxy02.pop.ord.plex.bz
137 pubsub.plex.bz
138 pubsub.plex.tv
139 raw.githubusercontent.com
140 redirector.googlevideo.com
141 res.cloudinary.com
142 s.gateway.messenger.live.com
143 s.marketwatch.com
144 s.youtube.com
145 s.ytimg.com
146 s1.wp.com
147 s2.youtube.com
148 s3.amazonaws.com
149 sa.symcb.com
150 secure.avangate.com
151 secure.brightcove.com
152 secure.surveymonkey.com
153 services.sonarr.tv
154 skyhook.sonarr.tv
155 spclient.wg.spotify.com
156 ssl.p.jwpcdn.com
157 staging.plex.tv
158 status.plex.tv
159 t.co
160 t0.ssl.ak.dynamic.tiles.virtualearth.net
161 t0.ssl.ak.tiles.virtualearth.net
162 tawk.to
163 tedcdn.com
164 themoviedb.com
165 thetvdb.com
166 tinyurl.com
167 title.auth.xboxlive.com
168 title.mgt.xboxlive.com
169 traffic.libsyn.com
170 tvdb2.plex.tv
171 tvthemes.plexapp.com
172 twimg.com
173 ui.skype.com
174 video-stats.l.google.com
175 videos.vidible.tv
176 vidtech.cbsinteractive.com
177 weather-analytics-events.apple.com
178 widget-cdn.rpxnow.com
179 win10.ipv6.microsoft.com
180 wp.com
181 ws.audioscrobbler.com
182 www.dataplicity.com
183 www.googleapis.com
184 www.msftconnecttest.com
185 www.msftncsi.com
186 www.no-ip.com
187 www.youtube-nocookie.com
188 xbox.ipv6.microsoft.com
189 xboxexperiencesprod.experimentation.xboxlive.com
190 xflight.xboxlive.com
191 xkms.xboxlive.com
192 xsts.auth.xboxlive.com
193youtu.be
194 youtube-nocookie.com
195 yt3.ggpht.com
196 zee.cws.conviva.com
197
198 # YouTube family base domains so every subdomain (video stream servers on
199 # *.googlevideo.com, image hosts on *.ytimg.com, profile thumbnails on
200 # *.ggpht.com, etc.) is automatically allowlisted via the AdGuard '||domain~'
201 # wildcard. The specific s.youtube.com / manifest.googlevideo.com entries
202 # above are kept for clarity but will be optimized away by rule-compiler
203 # because their parents are now in this list.
204 youtube.com
205 googlevideo.com
206 ytimg.com
207 ggpht.com
208 gvt1.com

```

```
209 gvt2.com
210 gvt3.com
211 gvt1.net
212 gvt2.net
213 gvt3.net
214
215 edge-mqtt.facebook.com
216 gateway.instagram.com
217 i.instagram.com
218
219 facebook.com
220 fb.com
221 fbcdn.net
222 fbsbx.com
223 messenger.com
224 whatsapp.com
225 web.whatsapp.com
226 instagram.com
227 cdninstagram.com
228 static.whatsapp.net
229 media-bos5-1.cdn.whatsapp.net
230 media-yyz1-1.cdn.whatsapp.net
231 webtp.whatsapp.net
232 mmg.whatsapp.net
233 g.whatsapp.net
234 g-fallback.whatsapp.net
```