Nullexit: Unified Project Document

Omar Alaaeldein

July 24, 2026

Contents

1 Project Directory Tree

```
nullexit/
|-- .aiignore
|-- .claude
|   `-- settings.local.json
|-- .env.example
|-- .gitignore
|-- LICENSE
|-- README.md
|-- Recover-Gateway.applescript
|-- Sweep-Gateway.applescript
|-- Toggle-Gateway.applescript
|-- agent.md
|-- benchmarks
|   |-- benchmark.sh
|   `-- test_load.py
|-- black_list.txt
|-- cloud-init.yaml
|-- devref.md
|-- diagrams.md
|-- docker
|   `-- tor
|       |-- Dockerfile
|       `-- entrypoint.sh
|-- docker-compose.yml
|-- launchd
|   |-- com.nullexit.noise.plist
|   |-- com.nullexit.taildrop-watch.plist
|   `-- com.nullexit.wake-recovery.plist
|-- post-rules.txt
|-- recover.sh
|-- scripts
|   |-- Dockerfile.routing-fix
|   |-- Dockerfile.rule-compiler
|   |-- Toggle-Gateway.bat
|   |-- check_circuits.py
|   |-- common.sh
|   |-- crypto.sh
|   |-- device-scramble.sh
|   |-- diagnose-host-leak.sh
|   |-- dns-proxy.py
|   |-- dns_anomaly_detector.py
|   |-- fix-docker-bridge-collision.sh
|   |-- fix-ssh-delay.sh
|   |-- generate_tex.py
|   |-- logger
|   |   |-- go.mod
|   |   `-- main.go
|   |-- noise.sh
|   |-- pf.conf
|   |-- pihole-mode.sh
|   |-- profiles.sh
|   |-- reinstall-host-tailscale.sh
```

```
|   |-- routing-fix.sh
|   |-- rule-compiler
|   |   |-- go.mod
|   |   `-- main.go
|   |-- setup-common.sh
|   |-- setup-linux.sh
|   |-- sweep.sh
|   |-- taildrop-watch.sh
|   |-- unlock-files.sh
|   |-- watcher.sh
|   `-- wifi-rejoin.sh
|-- setup.sh
|-- toggle.sh
|-- tor-bridges.txt
`-- white_list.txt
```

2 README.md

```
# nullexit: Tailscale + Cloudflare WARP Docker Gateway

> **Last updated:** July 23, 2026  [Architecture & Flow Diagrams ](./diagrams.md)  [Full Dev Reference ](./devref.md)

**nullexit** is a chained network gateway that routes all Tailscale exit-node traffic through a Cloudflare WARP VPN tunnel  double-encrypting every packet, hiding your ISP metadata, and providing network-wide DNS ad-blocking (AdGuard Home) and kernel-level IP threat blocking (`ipset`/`iptables`) for every device on your mesh.

---

## 1. Prerequisites
- Docker and Docker Compose installed.
- A Tailscale account.
- **macOS:** Install the standalone Tailscale CLI via Homebrew (`brew install tailscale`) and register `tailscaled` as a system service (`brew services start tailscale`). No `.app` GUI install required.
- **OS & Python:** Developed and tested against **macOS 26.5.2**. The project has zero external Python dependencies (no `requirements.txt`) and explicitly targets the default Apple-provided **Python 3.9.6** (`/usr/bin/python3`). While the scripts could function on newer versions, `nullexit` explicitly avoids experimenting with or modifying the built-in system Python environment to prevent unexpected OS-level side effects.

## 2. Installation & Setup

Run the interactive setup script for your OS. It downloads `wgcf`, generates fresh Cloudflare WARP WireGuard keys, prompts for a Tailscale Auth Key, and writes your `.env`.

```bash
./setup.sh          # macOS
./scripts/setup-linux.sh  # Linux
```

Reference `.env`
This is a quick reference of the common variables. For the **complete, authoritative list** of every `.env` variable (including all optional/advanced settings), see `devref.md 7`.
```env
WIREGUARD_PRIVATE_KEY=<Generated>
WIREGUARD_PUBLIC_KEY=<Generated>
WIREGUARD_ADDRESSES=<Generated>
TS_AUTHKEY=tskey-auth-...
GATEWAY_RULE_PROFILE=medium        # light | medium | heavy
# Safe MSS for double-tunneled traffic (Tailscale + WARP). Don't raise this
# devref.md 15.3.2 found 1180 (and even 1160) still too large and causes the
# exact mysterious mid-transfer stalls this value exists to prevent.
GATEWAY_MSS=1120
# Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own internet is NOT hijacked).
GATEWAY_HIJACK_HOST=true
GATEWAY_USE_EXIT_NODE=true
WARP_FAIL_THRESHOLD=6              # consecutive warp=off polls before auto-shutdown (default 6 = 30s)
KILL_SWITCH=true                   # enforce strict PF lock that breaks SSH if VPN fails. Leave true  this is the core leak guarantee.
# On campus/enterprise networks (WPA2-Enterprise / 802.1x), AP Client Isolation blocks direct
# LAN traffic between devices. Tailscale attempts a local P2P upgrade anyway, which causes
# macOS gvproxy to SNAT-mangle the reply's source port  poisoning the phone's WireGuard
# endpoint and blackholing 100% of its traffic. Setting this to false drops those local
# hole-punch packets so the phone safely falls back to Tailscale DERP relays.
# watcher.sh auto-detects 802.1x and AP isolation on every Wi-Fi change and overrides
# this setting automatically  you only need to set this manually if auto-detection fails.
TAILSCALE_ALLOW_LAN_P2P=false     # auto-managed; true only on trusted hotspots/home networks
ADGUARD_USER=admin                 # Username for AdGuard Home (defaults to admin if not set)
ADGUARD_PASSWORD=nullexit          # Shared password for AdGuard Home and Tor ControlPort (defaults to nullexit if not set)
BLOCKED_COUNTRIES="kp il"          # dynamically block country IP ranges via ipdeny.com
TOR_EXCLUDE_EXIT_NODES={us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il}  # comma-separated country codes to exclude as Tor exit nodes
DEBUG_TRACE=false                  # true = mirror a full command-by-command trace of the toggle to output.log (see 9)
```

3. Deploy the Gateway

macOS (Recommended):
```bash
./toggle.sh
```
```

```
Running it again stops the gateway and restores your network. Handles the full lifecycle: Colima VM, containers, DNS hijacking, Tailscale exit-node routing, rule compilation, and sleep prevention.

**Linux / Manual:**
```bash
docker compose up -d
```

4. Post-Deployment

Approve the Exit Node
1. Go to [Tailscale Admin Machines](https://login.tailscale.com/admin/machines).
2. Find the `tailscale` device `...` **Edit route settings** enable as Exit Node.

Verify Traffic Flow
On any Tailscale client device, select this gateway as the exit node, then visit api.ipify.org or run:
```bash
curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
# expect: warp=on
```

AdGuard DNS Setup

To ensure client devices correctly route DNS through the AdGuard container (and do not bypass it via `tailscaled`'s internal resolver), you **must** configure your Tailscale Admin Console:
1. Go to the **DNS** tab in the [Tailscale Admin Console](https://login.tailscale.com/admin/dns).
2. Enable **MagicDNS**.
3. Under **Nameservers**, add a **Custom** nameserver and enter the gateway's Tailscale IPv4 address (run `tailscale ip -4` on the gateway machine to find it, e.g., `100.x.x.x`).
4. Click the **three dots (...)** next to the nameserver you just added and enable **"Use with exit node"**.
5. Delete any other Global Nameservers (like Google or Cloudflare).
6. Toggle ON **Override local DNS**.

Traffic will now be correctly forced into the AdGuard sinkhole, and mobile devices will be blocked from bypassing it via encrypted DNS (DoH/DoT).

> [!WARNING]
> **Disable "Private DNS" on mobile devices.** Android's Private DNS (DoT/DoH) bypasses port 53 interception entirely. Go to `Settings  Connections  More connection settings  Private DNS  Off`.

> [!WARNING]
> **Bandwidth doubling:** Streaming 1GB of video on a remote device consumes 1GB download + 1GB upload on the host's network. Be mindful on metered connections.

AdGuard Dashboard
Navigate to `http://<gateway-tailscale-ip>:3000` credentials: **`admin` / `nullexit`**.

Access Any Mesh Device From Anywhere (SFTP/SSH)

Once the gateway exit node is online, **any device on your Tailscale mesh can reach any other mesh device, anywhere in the world** as long as both devices are on the mesh and the target device has an SSH server running. No port forwarding, no public IPs, no firewall holes. You can SSH into any device using its MagicDNS hostname (e.g. `ssh username@my-macbook`) or its assigned Tailscale IP (`100.x.x.x`).

> **Security Tip (Zero Local Attack Surface):** It is highly recommended to disable macOS's native "Remote Login" and "Screen Sharing" in System Settings to prevent exposing those ports (22, 5900) to physical Wi-Fi networks (which scanners can fingerprint). The gateway enforces `tailscale up --ssh=true` and the `pf` Kill-Switch automatically blocks local Wi-Fi access to these ports. Since both MagicDNS and Tailscale IPs route through the exact same encrypted tunnel, they are equally securebut **MagicDNS is recommended simply for ease of use.**

Extreme OPSEC: The Secret Tor-over-WARP Proxy
`nullexit` includes a completely isolated, headless Tor infrastructure proxy designed for terminals and hacking tools (e.g., `curl`, `sqlmap`, `nmap`).
To protect against local malware fingerprinting:
1. **Procedural Ports:** The proxy ports are randomized into the ephemeral range (10000-65000) on every single boot using a strict kernel socket collision-avoidance check (`netstat`).
2. **Honey-Port Tripwire:** The gateway spawns a fake trap port. If any malware attempts a sequential port-scan on `localhost` to find the proxy, it trips the Honey-Port, instantly triggering a macOS desktop notification and a critical log alert.

> [!NOTE]
> **Threat Model limitation:** Port randomization and the Honey-Port only defeat blind, generic port scanners. They do not protect against targeted local malware that runs `lsof -i` or reads the `/tmp/nullexit-ports.env` file directly. To actually prevent malicious local processes from reaching the proxy, you must run a host-level outbound firewall (like LuLu 4.3.0+ or Little Snitch) with **localhost/loopback filtering explicitly enabled** to gate access by process identity.

To use the Tor proxy, run `cat /tmp/nullexit-ports.env` to find your `$TOR_SOCKS_PORT` for the current session, and point your hardened browser (LibreWolf/Tor Browser) or CLI tool to `127.0.0.1:$TOR_SOCKS_PORT`.
```

```
119
120 > [!WARNING]
121 > **Preventing DNS Leaks (SOCKS5 vs. SOCKS5-Hostname):**
122 > When routing command-line tools through Tor, you **must** force the tool to delegate DNS resolution to the
      Tor proxy. Using plain SOCKS5 (e.g., `socks5://` or curl's `--socks5` flag) will resolve hostnames *locally
      * using the host's DNS settings (AdGuard/WARP) before establishing the connection. While this DNS traffic
      is encrypted inside the WARP tunnel, the destination domain name is still leaked to AdGuard and Cloudflare,
       undermining your anonymity.
123 >
124 > Use these correct protocols and configurations:
125 > - **`curl`:** Use `--socks5-hostname` or the `socks5h://` scheme (e.g., `curl -x socks5h://127.0.0.1:
      $TOR_SOCKS_PORT https://check.torproject.org`). Do **not** use `--socks5` or `socks5://`.
126 > - **`sqlmap`:** Pass the proxy URL using the `socks5h://` scheme to ensure remote resolution: `--proxy="
      socks5h://127.0.0.1:$TOR_SOCKS_PORT"`.
127 > - **`nmap`:** Nmap's built-in `--proxies` option resolves hostnames locally. To safely scan targets without
      DNS leaks, tunnel it using `proxychains-ng` configured with `proxy_dns` enabled (add `socks5 127.0.0.1 <
      PORT>` to `/etc/proxychains.conf` or a local config, then run `proxychains4 nmap <args>`).
128
129
130 #### Transparent .onion Browsing & Dark Web Access
131 In addition to the randomized SOCKS5 proxy port, `nullexit` provides seamless, transparent `.onion` domain
      browsing for **all devices** on your Tailscale mesh:
132 - **Zero-Config DNS:** AdGuard Home automatically intercepts queries for `.onion` domains and resolves them to
      Tor's virtual mapping subnet.
133 - **Tailscale Auto-Routing:** The virtual subnet is mapped to the `198.18.0.0/15` benchmark range (RFC 2544).
      Since this range is not a private IP subnet, Tailscale's Exit Node mode automatically routes it to the
      gateway (bypassing local LAN exclusions).
134 - **Transparent NAT Proxying:** Kernel firewall rules in the gateway redirect all TCP traffic destined for
      `198.18.0.0/15` directly into Tor's transparent port (`9040`).
135 - **Exit Node Exclusion:** Supports the ability to exclude specific countries from being used as Tor exit nodes
       via the `TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`.
136
137 ##### Enable in Web Browsers
138 To prevent accidental DNS leaks, web browsers block `.onion` lookups by default. To browse dark web sites
      natively:
139 - **Firefox:** Go to `about:config`, search for `network.dns.blockDotOnion`, and set it to `false`. Type any
      onion address (e.g. `duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion`) directly into the URL
       bar and it will load.
140 - **Mobile Browsers (iOS/Android):** Connect to your Tailscale exit node, open Firefox (with blockDotOnion
      disabled), and browse onion sites natively on the go.
141
142 #### Hardening: Using obfs4 Tor Bridges
143 If you want to disguise your Tor traffic from the WARP exit node provider, you can optionally enable Tor
      bridges:
144 1. Obtain private `obfs4` bridge lines from [bridges.torproject.org](https://bridges.torproject.org) or the
      official Telegram bot.
145 2. Create a file named `tor-bridges.txt` in the root of the `nullexit` folder and paste your bridge lines
      inside it (one per line).
146 3. Set `TOR_USE_BRIDGES=true` in your `.env` file and restart the gateway. The proxy will refuse to start if
      the bridges file is empty or missing.
147
148 #### Example: Browse your Android phone's files from your Mac
149
150
151 1. **On your Android phone**, install [Termux](https://termux.dev/) and [Termux:Boot](https://f-droid.org/
      packages/com.termux.boot/) from F-Droid.
152 2. In Termux, install and start the SSH server:
153    ```bash
154    pkg install openssh
155    sshd
156    ```
157    Termux defaults to **port 8022** (ports below 1024 require root on Android).
158 3. Find your phone's Tailscale IP:
159    ```bash
160    tailscale ip -4   # e.g. 100.87.42.87
161    ```
162 4. **Stop Android from killing Termux**  do ALL of these (Samsung One UI is especially aggressive):
163    - **Settings  Apps  Termux  Battery  Unrestricted**
164    - **Device Care  Battery  Background usage limits  Sleeping apps / Deep sleeping apps**  remove Termux if
      listed
165    - In Termux, run `termux-wake-lock` to hold a persistent wakelock
166 5. **On your Mac/PC**, open [Cyberduck](https://cyberduck.io/) (or any SFTP client: WinSCP, FileZilla, FE File
      Explorer on iOS) and connect to:
167    - **Server:** `100.87.42.87` (your phone's Tailscale IP)
168    - **Port:** `8022`
169    - **Protocol:** SFTP (SSH File Transfer Protocol)
170    - **Username:** (your Termux username, usually `u0_a123`  run `whoami` in Termux to check)
171    - **Password:** (your Termux password  set with `passwd` in Termux if you haven't already)
172
173 That's it  you can now browse, download, and upload files to your phone from any device on the mesh, no matter
      where either device is physically located. This works identically for any mesh device: Mac  Mac, Windows
```

```
    Linux, phone  NAS, etc.

> **Troubleshooting:** If the connection hangs for ~30 seconds before failing, your phone's SSH server isn't running. Open Termux and run `sshd` again. If Android keeps killing it despite Unrestricted battery, the `termux-wake-lock` command is your fix.

---

## 5. Multi-Layer Firewall & Kill-Switch

### Layer 1  DNS Sinkhole (AdGuard Home)
Blocks ads and tracking domains before they are downloaded. In automated Lighthouse tests on ad-heavy sites:
- **Time to Interactive:** ~22% faster (51.0s  41.8s)
- **Speed Index:** ~20% faster (15.3s  12.8s)

Customize blocking via `black_list.txt` and `white_list.txt`. Rule profiles (`light` / `medium` / `heavy`) are set in `.env`.

### Layer 2  Kernel IP Firewall (`ipset`/`iptables`)
On every startup, the `rule-compiler` fetches threat-intelligence feeds (Spamhaus DROP, Feodo Tracker, Emerging Threats, CINS) and compiles **~16,700 unique malicious IPs/CIDRs** into the kernel `FORWARD` chain blocking both outbound C2 connections and inbound attack traffic. Zero configuration required.

### Layer 3  Geo-IP Blocking
Dynamically blocks all traffic to and from specific countries using live IP ranges from `ipdeny.com`.
To add countries to your blocklist, open your `.env` file and set the `BLOCKED_COUNTRIES` variable with 2-letter ISO country codes (e.g. `BLOCKED_COUNTRIES="kp il cn ru"`), then restart the gateway.

### Layer 4  macOS PF Kill-Switch
A native macOS Packet Filter (`pf`) kill-switch that enforces a strict default-deny policy on your physical Wi-Fi interface (`en0`). When `KILL_SWITCH=true` is set in your `.env`, it drops all outgoing traffic except the Cloudflare WARP endpoints and Tailscale DERP relays.

- **Failsafe Design:** If the VPN tunnel crashes or Docker dies, your Mac completely loses internet access rather than leaking your real IP.
- **Remote-Access Safe:** Carefully designed to whitelist Tailscale's control plane (`192.200.0.0/24`), `utun*` tunnel traffic, and local LAN subnets. This guarantees that your remote SSH sessions and local AirDrop will survive even when the kill switch engages.
- **Setup Requirement:** The toggle scripts need permission to manipulate the network and firewall in the background without prompting you for a password. You must configure your passwordless sudo config by running exactly:
  ```bash
  echo "$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /usr/bin/dscacheutil, /usr/bin/killall, /usr/bin/pkill, /bin/kill, /sbin/route, /sbin/ifconfig, /usr/bin/true, /opt/homebrew/bin/brew, /usr/local/bin/brew, /usr/bin/python3 -I $PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I $PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I $PWD/scripts/dns-proxy.py" | sudo tee /etc/sudoers.d/nullexit
  ```

---

6. Toggle Scripts

macOS `toggle.sh`
Fully automated lifecycle script. Also supports a native `.app` launcher:
```bash
osacompile -o "Toggle Gateway.app" Toggle-Gateway.applescript
```

Optional `.env` settings (common subset; the full canonical table lives in `devref.md 7`):
- `GATEWAY_BYPASS_PING=true` proceed even if pre-flight connectivity checks fail.
- `GATEWAY_USE_EXIT_NODE=false` DNS-only mode, skip exit-node routing.
- `GATEWAY_HIJACK_HOST=false` skip DNS hijacking on the host (provides VPN/adblocking only to external Tailscale peers).
- `GATEWAY_MSS` override the TCP MSS clamp for the double-tunnel. Default (and only validated-safe value) is `1120`; devref.md 15.3.2 found 1180 and even 1160 still too large for the Tailscale+WARP double-encapsulation overhead, causing the exact mid-transfer stalls this clamp exists to prevent. Applies to both the container (`post-rules.txt`) and the host (`pf.conf`) they're kept in sync automatically.
- `HOST_MTU=1200` override the host Tailscale interface MTU. Defaults to 1200 to fit inside the 1280-byte WARP tunnel.
- `KILL_SWITCH=true` enforce strict network lock that breaks SSH if the VPN fails.
- `STOP_COLIMA_ON_EXIT=true` fully shut down the Colima VM on toggle-off (saves battery on dedicated hosts).
- `TAILSCALE_ALLOW_LAN_P2P=true` allow direct Tailscale LAN P2P connections between devices on the same network.
  - **Default: `false`.** On campus or enterprise Wi-Fi (WPA2-Enterprise / 802.1x), AP Client Isolation blocks local device-to-device traffic. When Tailscale still attempts a local P2P upgrade, macOS's `gvproxy` layer SNAT-mangles the reply's source port. WireGuard's endpoint roaming treats this as a legitimate address change and updates the phone's outbound path to the newly-mangled port which has no listener. This blackholes 100% of the phone's traffic while the DERP relay path is abandoned. Setting this to `false` preemptively drops all RFC1918-destined Tailscale UDP from the gateway so the phone receives a clean failure and safely falls back to DERP.
```

```
  - **Auto-managed:** `watcher.sh` automatically detects WPA2-Enterprise via `system_profiler SPAirPortDataType` (the `airport` binary was removed in macOS 14.4) and probes for AP isolation by pinging the default gateway (`route -n get 1.1.1.1`) on every network change. It writes the detected value to `.lan_p2p_detected` in the repo root, which `routing-fix.sh` re-reads every 30 seconds  **no restart required.** Only set this manually if auto-detection fails. See `devref.md 15.7.1` for the full SNAT endpoint poisoning analysis and `devref.md 15.11.6` for the `airport` removal background. *(Note: Direct P2P between the gateway container and the host Mac itself is always permitted via the Docker bridge to bypass DERP relays, regardless of this setting).*
  - Set to `true` on trusted home networks or Windows/mobile hotspots (no AP isolation) for a faster, lower-latency direct path.
  - **The real payoff  untrusted-but-open networks (hotels, cafes, airports).** These aren't enterprise/802.1x networks, so they typically don't AP-isolate  `watcher.sh`'s auto-detection usually finds P2P viable and flips this on by itself. That's two kinds of protection at once with zero manual setup on that specific network: every byte between your mesh devices is WireGuard-encrypted whether it goes direct or via DERP, so another guest on the same open hotel Wi-Fi who can technically reach your device at Layer 2 still can't read or tamper with anything; and since you're already Tailscale-enrolled, none of this needs configuring per-network  no VPN profile to import, no setting to flip at each new hotspot. The same phone that was safely DERP-only on an isolated campus network this afternoon gets fast, secure P2P at the hotel tonight automatically, because `watcher.sh` re-probes isolation on every Wi-Fi change.
- `WARP_FAIL_THRESHOLD=3`  number of consecutive failed checks before forcing a gateway teardown (default: 6 checks = 30s).
- `WARP_ENDPOINT_1` / `WARP_ENDPOINT_2`  override the Cloudflare WARP WireGuard endpoints.

### Linux  `toggle.sh` / `recover.sh`
```bash
./toggle.sh    # toggle ON/OFF   (cross-platform, dispatches on $OSTYPE)
./recover.sh   # recovery tool   (cross-platform)
```

Sleep Prevention (macOS)
When the gateway is ON, `caffeinate -i -s` runs in the background: `-i` blocks idle system sleep (even on battery), and `-s` additionally blocks sleep outright whenever the Mac is on AC power (a no-op on battery, so unplugged runs still sleep normally). The display can still sleep normally. Without `-s`, real idle-sleep cycles were still slipping through and dropping every mesh device's exit node until the next wake-recovery cycle see `devref.md 15.5.8`.

Auto-Recovery (macOS post-wake / post-roam)
Install the launchd LaunchAgent so the gateway self-heals after sleep/wake and Wi-Fi roams:
```bash
cp ./launchd/com.nullexit.wake-recovery.plist ~/Library/LaunchAgents/
sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```
Wi-Fi roams and network changes are caught reliably. Lid-close/wake is
best-effort on macOS Sonoma+ Apple suspends LaunchAgents during forced
sleep rather than re-triggering them on every wake, so the watcher can miss
the specific wake event (other self-healing the DNS Watcher, Tailscale's
DERP fallback, gluetun's healthcheck still recovers within ~30-60s
regardless). If lid-wake responsiveness matters more than that window,
`brew install sleepwatcher` and hook `recover.sh --post-wake` to `~/.wake`
instead. See `devref.md 15.5.1` for the full deep dive.

Wi-Fi auto-rejoin (macOS)
If this Mac drops Wi-Fi and macOS **auto-join is off**, it will otherwise just sit
there with no internet (and, if the gateway is up, the WARP Watcher would
eventually mistake the dead uplink for a dead tunnel see `devref.md 15.6.2`).
The same always-on `watcher.sh` daemon now also re-associates your Wi-Fi for you.

It never scans for or hops onto arbitrary networks macOS hides the live SSID
from every script unless you grant Location (which this project deliberately does
not use; see `devref.md 15.11.10`), so instead you **remember** your network
name(s) once and it only ever rejoins one of those:
```bash
bash scripts/wifi-rejoin.sh set  "Your-Wi-Fi-Name"   # remember a network
bash scripts/wifi-rejoin.sh list               # show remembered networks
bash scripts/wifi-rejoin.sh forget "Your-Wi-Fi-Name"  # stop remembering it
```
The remembered list lives in `.last_wifi_ssid` (one per line, gitignored). When
the Wi-Fi interface has no routable IP, `watcher.sh` calls
`scripts/wifi-rejoin.sh rejoin`, which waits a short grace period (let DHCP
self-heal first), then re-associates a remembered network via
`networksetup -setairportnetwork` pulling the password from your Keychain with
progressive backoff and a give-up cap so it never churns a network that won't come
up. It is leak-safe: it only acts when the interface has **no** address, i.e. no
egress path exists. Because the driver is the always-on daemon (not the gateway's
WARP Watcher), it reconnects you whether the gateway is up, down, or off. See
`devref.md 15.5.9`.

Config profiles taming the flag sprawl
The opt-in features below add up to several interacting `.env` switches. Rather than
hand-tune them, `scripts/profiles.sh` bundles the intent-level flags into three
```

````
known-good presets. It is **config-only**  it just edits `.env`, never touches the
live gateway  so changes apply on your next `./toggle.sh --restart`, and it forces
`KILL_SWITCH=true` in every profile.

```bash
bash scripts/profiles.sh show           # the whole current config, grouped + explained
bash scripts/profiles.sh preview home   # exactly what 'apply' would change (read-only)
bash scripts/profiles.sh apply home     # write it (backs up .env; never restarts)
```

Profiles: **campus** (AP-isolated enterprise, conservative)  **home** (trusted, fast
direct P2P, FaceTime + LAN Pi-hole)  **hardened** (cover traffic + Tor bridges, no
real-IP exposure).

### Cover traffic / dummy padding (macOS  optional, opt-in)
`scripts/noise.sh` emits verifiable-random UDP padding toward the exit so a passive
on-path observer sees a continuously varying encrypted flow instead of your real
traffic envelope (a metadata / traffic-flow correlation defence). Unlike the rest
of the toolkit it is **active**  it puts traffic on the wire  so it stays inert
unless you opt in. **One switch controls everything:** set `NOISE_ENABLED=true` in
`.env`. There is no separate "persistent vs one-shot" option  installing the
LaunchAgent below simply means padding auto-starts whenever `NOISE_ENABLED=true`,
and the job is a clean no-op whenever it is `false`.

```bash
# One-off (this session only):
bash scripts/noise.sh verify   # prove the noise is statistically random (entropy/monobit/chi-square)
bash scripts/noise.sh start    # arm padding now;  `status` / `stop` to manage

# Persistent (survives reboot/login)  install the LaunchAgent once:
cp ./launchd/com.nullexit.noise.plist ~/Library/LaunchAgents/
sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" ~/Library/LaunchAgents/com.nullexit.noise.plist
launchctl load -w ~/Library/LaunchAgents/com.nullexit.noise.plist
# To stop persisting:  launchctl bootout gui/$UID/com.nullexit.noise
```
Tunables (rate, packet size, jitter, target) live in `.env` under the
`NOISE_PAD_*` keys; defaults are a modest 64 kbps toward the gateway. Honest
limits: this is one-way random padding  it blurs uplink volume + timing, not a
full traffic-analysis-resistant shaper. See the `noise.sh` / `Cover Traffic &
Padding` notes in the knowledge graph.

### LAN Pi-hole mode (macOS  optional, opt-in)
nullexit is already a Pi-hole for **mesh** devices  AdGuard Home filters DNS for
everything that routes through the gateway over Tailscale. This mode extends that
filtering to **non-Tailscale** devices on the same LAN (a smart TV, a guest, IoT):
point their DNS server at this Mac's IP and they get the same ad/tracker/threat
blocking. It reuses `dns-proxy.py` bound to the LAN interface, forwarding to
AdGuard.

```bash
# .env: PIHOLE_LAN_MODE=true   (default false)
sudo bash scripts/pihole-mode.sh enable    # bind the LAN DNS forwarder (:53 needs root)
bash scripts/pihole-mode.sh test           # prove it resolves clean + sinkholes blocked
bash scripts/pihole-mode.sh status         # running? listen addr, AdGuard reachability
sudo bash scripts/pihole-mode.sh disable
```

Honest limits: it only works on a **trusted, non-isolated** network  on
AP-isolated WPA2-Enterprise (the same isolation that forces phones onto DERP)
other devices can't reach this Mac at all. And it's **DNS filtering only**  it
does not route those devices through WARP. See the `pihole-mode.sh` note in the
knowledge graph.

### Device Scramble Mode (macOS  optional, opt-in)
Randomizes the identifiers a network/observer uses to fingerprint your device (hostname and MAC address). This is useful on open / no-auth Wi-Fi (caf, hotel, airport) where these IDs are the only handle on you.
* Note: MAC spoofing requires root. The `test-mac` command is fail-safe and detects Apple-Silicon spoof-blocking without stranding you.

```bash
bash scripts/device-scramble.sh status            # show current + saved identity
sudo bash scripts/device-scramble.sh scramble     # random hostname (add --mac to also rotate MAC)
sudo bash scripts/device-scramble.sh restore      # put your real identity back
sudo bash scripts/device-scramble.sh test-mac     # safe test to see if this network takes a random MAC
```

---

## 7. Networking Notes
````

- **Port conflicts:** None. Tailscale uses dynamic UDP ports; WARP uses outbound UDP:2408. The gateway binds `0.0.0.0:41642/udp` on the host for Tailscale's WireGuard listener this is intentional and required to receive inbound hole-punch packets from peers on the internet. All other internal ports (AdGuard :5335, SOCKS5 :1080) are bound to `127.0.0.1` only and are not reachable from the LAN.
- **Tailscale P2P vs DERP:** On the same Wi-Fi, Tailscale establishes a fast P2P connection. On cellular (CGNAT), it always falls back to a DERP relay (~80-200ms). See `devref.md 5.1`.
- **AirDrop / AirPlay:** Unaffected the gateway only touches standard Wi-Fi/Ethernet interfaces, not `awdl0`.
- **VPN coexistence:** Do **not** run a local VPN client (WARP, Mullvad, NordVPN) simultaneously with the exit node. Both fight for the default route and will blackhole your connection.
- **Banking & financial sites:** May intermittently block logins through WARP. Banks use anti-fraud databases that flag datacenter IP ranges (like Cloudflare's `104.28.x.x`) as VPN/proxy traffic. Success fluctuates depending on which WARP IP you land on. If a site blocks you, either disable the exit node temporarily for that session or use the bank's mobile app (which often uses different fraud detection).
- **MSS clamping (Bufferbloat):** Double-tunnelling adds ~120-160 bytes of overhead per packet. This is now automatically handled natively via the macOS `pf` firewall, which universally applies TCP MSS Clamping to all outbound packets to strictly prevent MTU fragmentation and bufferbloat. No manual `.env` configuration is required. Additionally, the host's Tailscale interface MTU is explicitly lowered to 1200 (configurable via `HOST_MTU` in `.env`) to guarantee packets fit inside the container's WARP tunnel without fragmenting.

---

8. Privacy & Security Hardening

The stack shifts trust across four layers: WARP hides traffic from your ISP, AdGuard blocks tracking at DNS, Tailscale encrypts device-to-device, and `ipset` blocks known malicious IPs at the kernel.

For higher security, the gateway supports these drop-in upgrades:

Layer	Default	Hardened
VPN Exit	Cloudflare WARP (US)	Mullvad (Swedish, no-logs, anonymous payment)
DNS Resolver	AdGuard via WARP	AdGuard via Mullvad DoH
Coordination	Tailscale (US)	Headscale (self-hosted)
Auth	Persistent OAuth keys	Ephemeral auth keys
Content	WireGuard asymmetric	WireGuard + Pre-Shared Keys (PSK)

Python Runtime Airgap (Isolated Mode)
Never install third-party `pip` packages (like `requests` or `pyyaml`) into the Python environment that `nullexit` uses. All python scripts in this project (e.g., the `dns-proxy.py` daemon) are strictly engineered to run on the zero-dependency Python Standard Library. To strictly isolate the environment from supply-chain attacks, `nullexit` executes these scripts using **Python's Isolated Mode** (`python3 -I`). This deliberately blinds the execution environment to any malicious user-installed `pip` packages on the host, permanently air-gapping the gateway from PyPI.

See `devref.md 9` for the full threat model, Snowden programme analysis, and trust assumptions.

---

9. Debugging

- **`output.log`** All stderr from `toggle.sh`, `recover.sh`, `setup.sh`, and the rule-compiler is written here. The WARP Watcher (`WARP_FAIL_THRESHOLD`) also writes its events here `WARP DOWN`, `WARP RECOVERED`, and `WARP SHUTDOWN` notices, plus `WARP watcher PAUSED/RESUMED` when it detects the physical uplink is down (15.6.2). The Wi-Fi auto-rejoin helper logs `wifi-rejoin:` lines (re-associate attempts, backoff, give-up). All timestamps are **UTC** (unified 2026-07-23 the log used to mix UTC and local time). Check this first.
  - Every run now records lifecycle breadcrumbs regardless of `DEBUG_TRACE`: an `EXEC:` / `EXIT <code>:` line brackets each heavyweight command (`colima start`, `docker compose up -d --build`, `docker compose down`), and an `EXIT: code=<N> phase=<init|stop|start>` line is written when the script exits **for any reason** including a killed terminal, `pkill`, or OOM. This closes a prior blind spot where an interrupted **START** (e.g. a `--restart` race while Wi-Fi was re-associating) left the log ending abruptly after the teardown with no indication of what failed. If a toggle "does nothing" or the gateway ends up half-up, `grep -E 'EXEC:|EXIT ' output.log | tail` shows the last command run and its result.
- **`DEBUG_TRACE=true`** (in `.env`) Full command-by-command execution trace of `toggle.sh`, mirrored to `output.log` for deep post-mortems of intermittent failures. Each traced line is timestamped and tagged with `file:line:function` (via a custom `PS4`), so you can follow the *exact* code path a run took including subshells, Docker/Colima calls, and where it aborted. Off by default because it is verbose (the log grows quickly); enable it only while reproducing a specific problem, then set it back to `false`. Implementation note: it prefers bash's `BASH_XTRACEFD` (bash 4.1, e.g. most Linux) to send the trace to the log while keeping your terminal clean; on macOS's stock `/bin/bash` 3.2 it transparently falls back to teeing stderr to the log (trace also appears in the terminal). It never uses the dynamic-fd (`exec {var}>>`) form, so it is safe on 3.2. To read a trace: `grep -nE '^\+ ' output.log`.
- **`bash scripts/diagnose-host-leak.sh`** One-shot host-routing diagnostic. Runs 8 checks and classifies your state into one of four known scenarios (System Extension conflict, SOCKS5 fallback, IPv6 leak, route-table freeze) or OK, and prints the exact fix command. Check 5 uses `route -n get 1.1.1.1` to ask the kernel which interface real traffic actually resolves through (more accurate than `netstat -rn | grep default` on macOS). Pass `--fix` to apply the remediation automatically. Pass `--watch` (or `--watch 30`) to run a full baseline diagnostic then continuously monitor warp/IPv6/default-route for leaks, alerting on any state change.
- **`bash scripts/sweep.sh`** One-shot, read-only gateway health sweep. Runs 7 checks containers up/healthy, double-tunnel/IP-leak (host egress == container egress with `warp=on`), direct P2P to the gateway (no DERP

```
     relay), AdGuard compiled-rule counts + live DNS interception, PF kill-switch armed, tailnet reachability (a
      *burst* of `tailscale ping`s to every online peer, reporting loss % and RTT jitter  not just one-shot
     latency  and flagging DERP-relayed peers as throughput-degraded), and real throughput (host path vs. a WARP
     -only baseline via the container's SOCKS5 proxy, against a fast CDN  no exec/install inside the container)
      then rolls them into a PASS/WARN/FAIL verdict. Writes a timestamped `sweep-<UTC>.txt` report to the repo
     root (diffable across runs) and exits non-zero on FAIL, so it can gate launchd hooks or CI. Pass `--quiet`
     for just the verdict line. It never mutates routing, PF, DNS, or containers  safe to run against a live
     gateway at any time.
- **`bash scripts/fix-docker-bridge-collision.sh`**  One-shot fix for Docker bridge IP conflicts. If your local
      Wi-Fi router assigns you an IP in the `172.17.0.0/16` range, it will violently collide with Docker's
     default internal bridge, permanently killing your internet. This script detects the collision and auto-
     patches your Colima VM to use a safe subnet (e.g. `10.200.0.1/24`).
- **`devref.md`**  Complete architecture reference, routing deep-dives, and full resolved-issues log. Read this
      before modifying any routing or firewall logic.

> [!CAUTION]
> **Two infinite routing loops are possible.** (1) If a hotspot device providing internet to the gateway host
     is *also* using the gateway as its exit node, packets bounce forever between the two. (2) When the host's
     default route is redirected to `utun*`, WARP endpoint packets can loop back into the tunnel. Both are
     documented in `devref.md 15.2.1` and `15.2.2` and are automatically handled by `toggle.sh`.

---

## 10. Unified Project Document (Quine)

The project includes a `generate_tex.py` script that compiles the entire source code, configurations, and
     documentation into a single, unified LaTeX document (`nullexit_unified.tex` / `.pdf`).

Funnily enough, this document acts as a "project-level quine." It encapsulates the entirety of its own source
     code, its documentation, and the exact Python code required to generate itself. It serves as a self-
     contained, long-term archiving strategyproviding everything needed to reconstruct and understand the `
     nullexit` system from a single file.

---

## 11. Acknowledgements
- **[SyameimaruKoa](https://github.com/SyameimaruKoa):** Dual-stack TCP MSS clamping, `SIGHUP` state-tracking
     in the routing sidecar, and `TS_AUTH_ONCE` integration.

## 12. License
GNU Affero General Public License v3. See [LICENSE](./LICENSE).

## 13. Changelog

- **July 23, 2026**  Wi-Fi resilience + logging cleanup. (1) **Wi-Fi auto-rejoin**: if the Mac drops Wi-Fi with
      auto-join off it no longer just sits there  the always-on `watcher.sh` daemon re-associates a *remembered*
      network (`scripts/wifi-rejoin.sh`; no Location, no SSID scanning) with a grace period, backoff, and give-
     up cap. Driven by the always-on daemon, not the gateway's WARP Watcher, so it works whether the gateway is
     up, down, or off. See `README 6` and `devref.md 15.5.9`. (2) **WARP Watcher no longer false-kills the
     gateway on a Wi-Fi drop**: it checks the physical uplink before counting a `warp=off` as a tunnel failure
     and PAUSEs (leak-safe: no egress = nothing to leak) instead of nuking the gateway  root cause of the
     2026-07-22 outage. The first cut only checked for a missing IP; a **stale-DHCP-lease** case on enterprise
     Wi-Fi slipped through and re-triggered the kill (2026-07-23, a 2h41m outage), so the guard now also reads
     the NIC link-layer state (`ifconfig  status: active`)  "has an IP" is not "has egress." See `devref.md
     15.6.2` and `15.12.27`. (3) **Logging unified to UTC**  `output.log`, `blocked.log`, and the AdGuard query
     log previously mixed UTC and local (EDT); everything is now UTC. See `devref.md 15.12.26`.
- **July 18, 2026**  Fixed two real bugs surfaced by live sweep evidence: the QUIC/DoT `REJECT` rules in `post-
     rules.txt` were shadowed by an earlier blanket `ACCEPT` (0 packets ever matched  the intended Instagram/
     QUIC-fragmentation fix from `devref.md 15.3.3` was never actually enforced), now reordered ahead of it; and
      the host-side `pf.conf` MSS clamp was stuck at the pre-`15.3.2` value of `1160` while the container had
     already moved to `1120`  `scripts/common.sh` now syncs both from `GATEWAY_MSS` so they can't drift apart
     again. Also fixed `setup.sh`/`scripts/setup-common.sh` never generating `NULLEXIT_SEED` (every fresh
     install hard-failed its first `toggle.sh` on a missing-seed crypto error) and defaulting `KILL_SWITCH` to `
     false` for new installs  see `devref.md 15.12.23`. `scripts/sweep.sh` gained a 7th check (**throughput**,
     host path vs. a WARP-only baseline against a fast CDN)  see `devref.md 15.3.6` for why the target choice
     matters and `9` below for the check itself. Check 6 (tailnet reachability) was upgraded from a single ping
     to a **burst** (loss % + RTT jitter per peer)  a peer can be 100% reachable and still have real quality
     problems a single ping hides; caught live on an S24 over DERP (0% loss, 665ms jitter). Confirmed that
     jitter isn't gateway-fixable (mobile-radio power state + DERP's TCP-relay timing, verified via a Wi-Fi-vs-
     cellular comparison) and that forcing P2P through confirmed AP isolation isn't possible either (Layer 2
     block, below anything Tailscale operates at)  see `devref.md 16.4`.
- **July 17, 2026**  `scripts/sweep.sh` (the read-only automated health sweep added July 15, mirroring `sweep.
     md 1`) gained a 6th check: **tailnet reachability**  every Tailscale peer reported online is probed with `
     tailscale ping` (TSMP, so it works even on phones that drop ICMP) and rolled into the PASS/WARN/FAIL
     verdict. Documented in 9 above.
- **July 12, 2026**  Tor ControlPort dynamic password authentication (unified with `ADGUARD_PASSWORD`),
     enabling live circuit inspection via `/sweep`. See `devref.md 15.10.2`.
- **July 11, 2026**  Tor container hardening: `obfs4proxy` restored (base image  `debian:bookworm-slim`),
     robust bridge-file validation, and image pinning. See `devref.md 15.10.1`.
- **July 10, 2026**  Roaming/startup stabilization suite: fixed the Wi-Fi-roam host-IP leak (raw ISP `132.x`)
     caused by the WARP Watcher racing post-wake recovery, plus a batch of pre-flight/concurrency/kill-switch/
```

```
    Tailscale-flag bugs and the macOS SNAT endpoint-poisoning P2P blackhole. Full post-mortems in the Incident
    Log  see `devref.md 15.2.7`, `15.5.3``15.5.5`, `15.2.8``15.2.9`, `15.7.1`, `15.7.3``15.7.5`, `15.12.18`.
- **July 6, 2026**  Edge-case security and stability hardening: DNS Watcher interface independence, robust lock
    -file PID verification, fail-closed crypto integrity check, strict `iptables` pinning for tailscale0tun0,
    Docker port binding to `127.0.0.1` to prevent LAN exposure, routing verification via `netstat -nr`. See `
    devref.md 17`.
- **July 4, 2026**  Colima bridged networking, SOCKS5 proxy migrated natively to Tailscale, headless prompt
    crashes fixed, proxy bypass domains added to fix LAN P2P. Python dependencies completely eliminated: `
    logger` and `rule-compiler` ported to blazing-fast Go multi-stage Docker builds. Added `crypto.sh` to
    enforce cryptographic signature validation on all core bash scripts. Massively refactored `toggle.sh` and `
    recover.sh` to eliminate ~200 lines of duplicated code by migrating network and process logic to `scripts/
    common.sh`. Added a concurrency mutex to `toggle.sh`, resolved unbound `TS_AUTHKEY` check crashes on setup,
     and integrated Go validation to the CI pipeline.
- **July 3, 2026**  **Critical Security Updates:** Addressed the overnight silent IP leak and improved geo-
    blocking. Introduced `scripts/host-leak-probe.sh` (a sub-second host-egress leak detector) and a
    statistical threshold auto-shutdown watcher. Added `unlock-files.sh` to safely reset file permissions.
    Resolved numerous startup regressions and system bugs.
- **July 2, 2026**  Two infinite routing loops resolved (`devref.md 15.2.2`): host-VM tunnel loop (WARP bypass
    routes now via gateway IP + manual `utun*` redirection) and Docker Compose subnet takeover (`10.200.1.0/24`
     lock). `--accept-routes=true` now explicit on all `tailscale up --reset` calls.
- **July 2, 2026**  Auto-recovery daemon documented in 6 above (`scripts/watcher.sh` + launchd plist). See `
    devref.md 15.5.1`.
- **July 2, 2026**  Linux scripts (`scripts/toggle-linux.sh`, `recover-linux.sh`, `setup-linux.sh`) documented
    in 6.
- **July 1, 2026**  `recover.sh --post-wake` ships; wired to watcher daemon. Triggered on every sleep/wake or
    network-state change.
- **June 28, 2026**  8 internal scripts moved to `scripts/`. User-facing files (`toggle.sh`, `recover.sh`, `
    setup.sh`, `docker-compose.yml`) stay at repo root.
- **June 28, 2026**  All hardcoded install paths removed. `SCRIPT_DIR`-relative resolution and `
    __NULLEXIT_HOME__` placeholder used throughout.

### Cryptographic Integrity Verification

nullexit uses a built-in cryptographic integrity checker designed to provide tamper-*evidence* against
    accidental edits, logic bugs, or non-privileged malware. During setup, a 256-bit `NULLEXIT_SEED` is
    injected into your `.env` file. The core entrypoint scripts (`toggle.sh`, `recover.sh`, etc.) are hashed
    using HMAC-SHA256, and their signatures are stored in `.signatures`.

*(Note: This is not a strict boundary against an attacker who already has write access to the repo, as they
    could simply read the seed and re-sign the scripts. It is designed as a strict footgun-catcher).*

If any script is modified without authorization, the gateway will loudly fail to boot. If you make intentional
    edits to the bash scripts, you must re-sign them by running:
```bash
./scripts/crypto.sh --sign
```
```

3 agent.md

```
# nullexit  Detailed Agent Analysis

> Complete per-file breakdown of every function, constant, duplication, and bug found.
> **Note (July 2026):** `sync-rules.py` and `logger.py` have been fully ported to Go (`scripts/rule-compiler/
    main.go` and `scripts/logger/main.go`) to dramatically improve performance and reduce container footprint
    via multi-stage Alpine builds. The analysis below reflects their previous Python states.

##  Important Agent Instruction
- **NEVER use `cat << EOF` or terminal redirections to write or modify files.** You must always use your native
     file-editing tools (`replace_file_content` or `multi_replace_file_content`). Using `cat EOF` leads to
    duplicated text, broken markdown headers, and structural corruption (which previously destroyed the section
     numbering in `devref.md`).
- **Always run `bash scripts/crypto.sh --sign` after making any modifications to the core bash scripts.** This
    is required to update the cryptographic HMAC-SHA256 signatures; otherwise, the startup integrity checks
    will fail and block execution.
- **Always run `git diff` and print the changes to the user *before* requesting or executing any git commit
    command.** The user must be shown the diff so they can review and understand the changes. Based on this
    diff, formulate a highly precise, detailed, and descriptive commit message (explaining exactly what was
    changed and why) rather than a generic summary, and present it to the user alongside the diff.
- **Always run `git status` before `git diff` to identify all modified and untracked files, ensuring no files (
    such as new scripts or configuration assets) are missed before formulating commit diffs and staging.**
- **Always use the `--restart` flag when asked to restart the toggle (e.g., run `./toggle.sh --restart`).**
- **NEVER execute any `sudo` commands (especially `sudo route`, `sudo dscacheutil`, or any network-modifying
    commands) without explicitly explaining to the user what the command does and why it is needed FIRST. Do
    not assume implicit permission. Wait for the user's approval before running it.**
- **When fixing an error and using logs to debug, always show the user the reasoning and the specific logs that
     support this theory.**
```

```
- **When the user tells you to "push", this means read git diffs and prepare commit messages that correspond to these changes (do not ignore any changes). If the changes can be atomized into multiple commits for easier understanding, do so.**
- **When the user says "latex this", it means you should run `python3 scripts/generate_tex.py` to regenerate the unified LaTeX document.**
- **NEVER apply changes to running gateway containers directly via `docker compose up` or `docker compose restart`.** Recreating or restarting containers (especially the `warp` container which hosts the network namespace) breaks the shared network stack for all other services (`adguardhome`, `tailscale`, `routing-fix`), severing host connectivity and freezing macOS DNS. If you need to apply changes or rebuild containers, cleanly stop the gateway first using `./toggle.sh`, or trigger `recover.sh --post-wake` to rebuild and re-verify the network stack safely in the correct dependency order. Do NOT run `./toggle.sh --restart` or any network-rebuilding commands yourself if the execution could sever host network access; instead, explain the changes and ask the USER to run `./toggle.sh --restart` (or `./toggle.sh` stop and start) themselves to safely apply the modifications.

##  How to Verify Gateway is Working
Whenever modifications are made to the gateway scripts, routing, or containers, execute these verification checks in order:
1. **Container Health:** Run `docker compose ps` to ensure all containers (`warp`, `adguardhome`, `routing-fix`, `tailscale`) are `running` and `healthy`.
2. **DNS Hijacking Status:** Run `networksetup -getdnsservers "Wi-Fi"` (or appropriate network service) and ensure it is set exclusively to the gateway's Tailscale static IP (typically `100.100.21.8`).
3. **DNS Query Interception:** Run `dig @100.100.21.8 google.com` (using the gateway static IP from `.gateway_ip`) to ensure AdGuard Home is active, intercepting, and resolving queries.
4. **Double-Tunnel Egress:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp` to verify the egress is routed through Cloudflare WARP (`warp=on`).
5. **Host Leak Monitoring:** Run `tail -n 30 output.log` and verify the background watchers (DNS + WARP) report healthy status without `LEAK` warnings. For a deeper on-demand audit, run `bash scripts/diagnose-host-leak.sh` (add `--watch` to poll, `--fix` to remediate).

**Automated:** `bash scripts/sweep.sh` runs a read-only 7-check health sweep covering the above plus the PF kill-switch, tailnet-wide peer reachability, and real throughput, writes a timestamped `sweep-<UTC>.txt` report to the repo root, and exits non-zero on FAIL (`--quiet` for just the verdict). It never mutates routing/PF/DNS/containers, so it is safe against a live gateway. See its entry in Part 3.

---

## Part 1: macOS Main Scripts

### toggle.sh (2191 lines, unified macOS + Linux)

> **Refreshed 2026-07-15.** Post-unification (`$OSTYPE` switch at L20) + the July-2026 `common.sh` extraction, the tables below reflect the *current* file. Most former "toggle.sh functions" now live in `scripts/common.sh`; the host-leak-probe pair was removed. For the ordered execution flow see **"toggle.sh  Execution Map & Critical Path"** above.

**Functions still defined in `toggle.sh`**  several exist in *both* halves (Linux L63848 / macOS L8782191):

| Function | Linux L | macOS L | Purpose |
|---|---|---|---|
| `_log_exit_breadcrumb()` | 30 | 885 | Write lifecycle-phase breadcrumb on exit (post-mortem trail) |
| `get_free_port()` | 67 | 944 | Pick a random unused port (self- + kernel-collision checked) |
| `run_gui_cmd()` |  | 1033 | Run a command as the logged-in console GUI user |
| `write_gateway_active_marker()` |  | 1053 | Write `/tmp/nullexit-gateway-active.marker` (macOS only) |
| `clear_gateway_active_marker()` |  | 1059 | Remove active marker + `TUNNEL_FAILED_CLOSED.marker` |
| `start_sleep_prevention()` | 145 | 1091 | Start `caffeinate` w/ revert trap |
| `stop_sleep_prevention()` | 188 |  | Stop caffeinate via PID file |
| `stop_dns_watcher()` | 202 |  | Stop DNS watcher (macOS uses `stop_pidfile_daemon`) |
| `start_warp_watcher()` |  | 1164 | Background WARP liveness monitor  `recover.sh` after N fails |
| `stop_warp_watcher()` |  | 1241 | Stop WARP watcher via PID file |
| `cleanup_handler()` | 215 | 1258 | Fail-closed teardown trap (guarded by `SUCCESS_RUN`) |
| `restart_tailscaled_daemon()` | 327 | *(common)* | Restart tailscaled (Linux-half local copy; macOS uses common.sh) |
| `disconnect_tailscale_host()` | 337 | *(common)* | `tailscale down` w/ retry (Linux-half local copy; macOS uses common.sh) |
| `setup_exit_node_routing()` |  | 1430 | Override default route through the Tailscale utun interface |
| `cleanup_network_state()` | 354 | 1464 | Nuclear cleanup: proxies, DNS cache, routes, Wi-Fi, sharing services |

**Delegated to `scripts/common.sh`** (moved out of toggle.sh in the July-2026 refactor; line = common.sh):

| Function | common.sh L | Function | common.sh L |
|---|---|---|---|
| `run_with_timeout()` | 58 | `enable_socks_proxy()` | 769 |
| `is_gateway_active()` | 168 | `disable_socks_proxy()` | 803 |
| `gateway_state()` | 232 | `reset_dns()` | 824 |
| `get_active_service()` | 284 | `stop_local_dns_proxy()` | 898 |
| `restart_tailscaled_daemon()` | 338 | `start_local_dns_proxy()` | 909 |
| `disconnect_tailscale_host()` | 352 | `force_dns_to_gateway()` | 988 |
| `stop_pidfile_daemon()` | 393 | `add_warp_bypass_routes()` | 484 |
| `start_dns_watcher()` | 715 | `remove_warp_bypass_routes()` | 551 |
```

```
68
69 **Removed:** `start_host_leak_probe()` / `stop_host_leak_probe()` and the whole `HOST_LEAK_PROBE` prober
       subsystem  no longer defined or called anywhere in `toggle.sh`, `recover.sh`, or `common.sh`, and `
       HOST_LEAK_PROBE` is no longer an env var (all verified). The only surviving host-leak tool is the
       standalone, on-demand `scripts/diagnose-host-leak.sh` (there is **no** `scripts/host-leak-probe.sh`).
70
71 **Constants / anchors (current):**
72
73 | Constant | Value | Line |
74 |----------|-------|------|
75 | `LOG_FILE` | `$PWD/output.log` | L22 (Linux) / L863 (macOS) |
76 | `LOCK_FILE` | `/tmp/nullexit-toggle.lock` | L1020 |
77 | `PID_FILE` (caffeinate) | `/tmp/nullexit-caffeinate.pid` | L142 |
78 | `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L200 |
79 | `WARP_WATCHER_PID_FILE` | `/tmp/nullexit-warp-watcher.pid` | L1144 |
80 | Gateway active marker | `/tmp/nullexit-gateway-active.marker` | L1054 (write) |
81 | TUNNEL_FAILED_CLOSED marker | `<repo>/TUNNEL_FAILED_CLOSED.marker` | L1061 |
82 | `SOCKS_PROXY_PORT` | random per-toggle, `1080` fallback | L91/L107 (Linux)  L968/L986 (macOS) |
83 | DNS search domain | `ts.net` | L763 |
84 | `PATH` (via `setup_standard_path`, common.sh) | Homebrew + system bins | L1330 |
85 | WARP endpoints (via `get_warp_endpoint_1/2`, common.sh) | `162.159.192.1` / `162.159.193.1` | common.sh L335
       -336 |
86 | DNS fallback | `1.1.1.1` (in `reset_dns`, common.sh) | common.sh L824 |
87
88 ---
89
90 ### toggle.sh  Execution Map & Critical Path (verified 2026-07-15)
91
92 > Traced from the current unified `toggle.sh` (2191 lines). The file carries **two halves** selected by `
       $OSTYPE` at **L20** (`if [[ "$OSTYPE" != "darwin"* ]]` runs the self-contained **Linux path  L63848** then
       `exit 0` at ~L845; darwin falls through to the **macOS path  L8782191**). Line numbers below are the **
       macOS path** (the deployment target). The halves are *similar in intent but not identical*  see "Linux path
         where it diverges" below. `~Line` = current, approximate.
93
94 #### Entry & dispatch
95
96 ```
97 ./toggle.sh [ (no-arg) | --restart | restart ]
98      |
99      |-- verify HMAC signatures (crypto.sh --verify) --fail--> exit 1   (L857)
100     |-- rotate output.log if > 1GB
101     |-- arm teardown trap:  ERR/INT/TERM/HUP -> cleanup_handler        (L1317-1320)
102     |-- generate random ports (Tor / SOCKS / DNS)                      (L940)
103     |
104     |-- --restart --> gateway_state? running -> run STOP, then START   (L999)
105     |
106     '-- no-arg   --> gateway_state?                                    (L1535)
107                        running  -> STOP  branch  (L1538-1584)
108                        stopped  -> START branch  (L1590-2191)
109 ```
110
111 `gateway_state` is driven by `/tmp/nullexit-gateway-active.marker` (persisted intent), falling back to a live
       container probe.
112
113 #### START  synchronous critical path (all must complete)
114
115 | # | Step | ~Line | Effect if it fails / is interrupted |
116 |---|------|-------|-------------------------------------|
117 | 1 | `disconnect_tailscale_host` (clean slate) | 1595 | host left mesh |
118 | 2 | `colima_start_until_ready` (readiness poll, 15.8.7) | 1631 | no Docker VM  abort |
119 | 3 | honey-port reap (`pkill`) + arm (`nc -l`, disowned) | 1695 | tripwire missing (non-fatal) |
120 | 4 | rule-compile (hash-gated, OFF critical path 15.8.8) | 1708 | reuse prior compiled rules |
121 | 5 | build-gate routing-fix / tor images | 1743 | rebuild only on change |
122 | 6 | `docker compose up -d`  warp, tailscale, routing-fix, adguard, tor | 1758 | no gateway |
123 | 7 | wait tailscaled reachable | 1907 | host tailscale unusable |
124 | 8 | **host `tailscale set --exit-node=$TS_IP`** | 2046 | **host not routed via gateway  REAL-IP LEAK** |
125 | 9 | `add_warp_bypass_routes` + `setup_exit_node_routing` | 2048 | WARP endpoint unreachable / wrong route |
126 | 10 | **`force_dns_to_gateway`** (+ `start_local_dns_proxy`) | 2078 | host DNS not hijacked |
127 | 11 | `enable_socks_proxy` + verify SOCKSWARP | 2100 | SOCKS off |
128 | 12 | print **"STARTED in Ns"** (timing marker) | 2127 | cosmetic only  NOT the commit point |
129 | 13 | `start_sleep_prevention` (caffeinate) | 2130 | Mac may sleep |
130 | 14 | `start_dns_watcher` (re-hijack DNS every 30s) | 2134 | no roam DNS recovery |
131 | 15 | `start_warp_watcher` (warp liveness  recover.sh) | 2138 | no WARP self-heal |
132 | 16 | **`write_gateway_active_marker`** | 2146 | watcher believes gateway is DOWN |
133 | 17 | **`SUCCESS_RUN=true`**  disarms the teardown trap | 2164 | *(true commit point)* |
134 | 18 | print "You can close this terminal window now." | 2191 |  |
135
136 #### STOP  critical path
137
138 `disconnect_tailscale_host` (1543)  `docker compose down --remove-orphans` (1552)  `colima stop` (1558)  `
       reset_dns` (1568)  `stop_local_dns_proxy` (1571)  stop caffeinate (1574)  stop DNS watcher (1575)  `
```

```
     stop_warp_watcher` (1576)  `clear_gateway_active_marker` (1577)  "STOPPED in Ns" (1584).
139
140  #### Long-lived background processes spawned by START
141
142  | Process | Started | PID / reap | Role |
143  |---------|---------|------------|------|
144  | honey-port `nc -l localhost $PORT` (disowned) | L1695 | `pkill -f "nc -l localhost"` | port-scan tripwire |
145  | `caffeinate` | L2130 | `/tmp/nullexit-caffeinate.pid` | prevent system sleep |
146  | DNS watcher loop | L2134 | `/tmp/nullexit-dns-watcher.pid` | re-hijack DNS every 30s (Wi-Fi roam) |
147  | WARP watcher loop | L2138 | `/tmp/nullexit-warp-watcher.pid` | monitor warp; PAUSE if uplink down / link `
     inactive` (15.6.2 / 15.12.27), else fire recover.sh after N fails |
148  | local DNS proxy `dns-proxy.py` | L2089 | (killed by `stop_local_dns_proxy`) | UDP:53  gateway DNS fallback |
149
150  Separately, the **launchd `scripts/watcher.sh`** daemon (post-wake / post-roam recovery) runs independently of
     toggle and invokes `recover.sh --post-wake`; toggle does **not** spawn it.
151
152  #### `cleanup_handler`  the fail-closed teardown trap
153
154  Armed on `ERR/INT/TERM/HUP` (L1317-1320), guarded by `SUCCESS_RUN`. While `SUCCESS_RUN=false` it runs: `
     reset_dns`, `disconnect_tailscale_host`, `stop_local_dns_proxy`, stop watchers, `
     clear_gateway_active_marker`, `docker compose down` (+ `colima stop` on `ERR`). This is the **fail-closed
     guarantee**: a half-built gateway is torn down, not left leaking.
155
156  ####  Correctness findings (verified against source)
157
158  1. **"STARTED in Ns" (L2127) is NOT the commit point  `SUCCESS_RUN=true` (L2164) is.** The marker write (L2146)
        and watcher starts (L2130-2138) sit *between* them. **Any TERM/INT  including a wrapping tool's timeout
     in the L2127L2164 window fires `cleanup_handler`  a FULL teardown** (containers + colima + host tailscale),
      leaving the marker ABSENT and the host on its raw ISP link, *despite* the printed success. This is exactly
      the 2026-07-15 incident. (The Linux half has the same shape but a *much* narrower window: STARTED L814  `
     SUCCESS_RUN=true` L820, and it maintains no marker.)
159  2. **Never pipe `toggle.sh` through a reader** (`| tail`, `| grep`). The disowned children (honey-port,
     watchers, dns-proxy, caffeinate) inherit stdout, so the pipe never EOFs and the reader hangs long after
     toggle has exited  which then trips the wrapper's timeout  SIGTERM  finding #1. Run detached with a file
     redirect instead: `nohup ./toggle.sh --restart > /tmp/t.log 2>&1 &`, then poll the file.
160  3. **Host wiring is correctly on the synchronous critical path** (steps 8 & 10 precede "STARTED"), so a *
     completed* START cannot leave the host unrouted/leaking  the only way to that state is an interruption per
     finding #1.
161  4. **Namespace ordering:** step 6's `docker compose up` creates warp's network namespace that `tailscale`/`
     routing-fix`/`adguard` share (compose `depends_on: warp healthy` enforces order). Never `docker compose
     restart warp` on a live gateway  it detaches the shared netns (see Agent Instruction, L16).
162
163  #### Linux path  where it diverges (L63848)
164
165  The Linux half is intent-similar but **not** a line-for-line mirror of macOS. Verified differences:
166
167  | Aspect | macOS (L8782191) | Linux (L63848) |
168  |--------|------------------|----------------|
169  | VM layer | Colima (`colima_start_until_ready`, L1631) | none  native Docker |
170  | Host exit-node assert | `$TS_BIN set  --exit-node="$TS_IP"` (L2046) | `$TS_BIN up --reset  --exit-node="
     $TS_IP"` (L735) |
171  | gateway-active marker | written (L2146) | **not maintained** (L119)  `gateway_state` degrades to a live
      container probe |
172  | Commit window | STARTED L2127  `SUCCESS_RUN` L2164 | STARTED L814  `SUCCESS_RUN` L820 |
173  | Port gen tooling | `jot` / `netstat` | `shuf`/`$RANDOM` / `ss` |
174
175  The `up --reset` vs `set` distinction matters: `--reset` re-applies *all* prefs and forces a route re-eval,
     whereas `set` mutates only the exit-node pref  the same asymmetry `recover.sh --post-wake` relies on (
     devref 15.12.21).
176
177  ---
178
179  ### recover.sh (566 lines, 23KB)
180
181  **Functions (7 total):**
182
183  | # | Function | Lines | Purpose |
184  |---|----------|-------|---------|
185  | 1 | `step()` | L40 | Print bold step header |
186  | 2 | `ok()` | L41 | Print green success message |
187  | 3 | `warn()` | L42 | Print yellow warning message |
188  | 4 | `fail()` | L43 | Print red failure message |
189  | 5 | `die()` | L44 | Print red error and exit |
190  | 6 | `run_with_timeout()` | L5674 | Run command with timeout (bash-native, no GNU coreutils) |
191  | 7 | `restart_tailscaled_daemon()` | L118129 | Restart tailscaled daemon via brew services |
192
193  **Constants:**
194
195  | Constant | Value | Line |
196  |----------|-------|------|
197  | Color codes | `RED`, `GREEN`, `YELLOW`, `BOLD`, `NC` | L37-38 |
```

```
| `SCRIPT_DIR` | `$(dirname "${BASH_SOURCE[0]}")` | L50 |
| `PATH` export | `/opt/homebrew/bin:/usr/local/bin:$PATH` | L77 |
| `PID_FILE` | `/tmp/nullexit-caffeinate.pid` | L387 |
| `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L407 |
| WARP endpoints | `162.159.192.1`, `162.159.193.1` | L329-336 |
| TUNNEL_FAILED_CLOSED marker | `$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker` | L51 |

**Duplicated logic from toggle.sh:**
*(Note: As of July 2026, **ALL** of the below duplicated logic has been successfully extracted to `scripts/common.sh`!)*
- `run_with_timeout()`  simpler subshell version vs toggle's PID-tracking version
- `restart_tailscaled_daemon()`  near-identical (uses warn()/ok() vs echo)
- Active network service detection  inline at L217-233 (toggle has `get_active_service()` function)
- EN0 service resolution  inline at L230-233 (same as toggle L515-516)
- WARP bypass routes  inline at L329-337 (toggle has `add_warp_bypass_routes()`)
- Exit node routing setup  inline at L340-345 (toggle has `setup_exit_node_routing()`)
- DNS cache flushing  L296-302 (same as toggle L611-616)
- Wi-Fi power-cycling  L442-461 (similar to toggle L626-637)
- Proxy disabling  L282-288 (same as toggle L604-608)
- Sharing services reset  L586 (same as toggle L642)
- PID file stop pattern  L387-399, L407-419, L423-435 (same as toggle L157-167, L193-203, L312-322)
- KILL_SWITCH check  L194 (same as toggle L141, L469)
- .gateway_ip reading  L138-142, L209-213 (same pattern as toggle L1080)

---

### setup.sh (410 lines, 18KB)

**Functions (4 total):**

| # | Function | Lines | Purpose |
|---|----------|-------|---------|
| 1 | `step()` | L10 | Print bold blue step header |
| 2 | `ok()` | L11 | Print green success message |
| 3 | `warn()` | L12 | Print yellow warning message |
| 4 | `die()` | L13 | Print red error and exit |

**Constants:**

| Constant | Value | Line |
|----------|-------|------|
| Color codes | `RED`, `GREEN`, `YELLOW`, `BLUE`, `BOLD`, `NC` | L7-8 |
| `RECOMMENDED_TS_VERSION` | `1.98.5` | L71 |
| `ADGUARD_PASSWORD` | `nullexit` | L215 |
| Default .env values | `GATEWAY_RULE_PROFILE=medium`, `GATEWAY_MSS=1120`, etc. | L240-248 |

---

## Part 2: Linux Scripts

### toggle.sh  Linux path (unified into the root cross-platform script)

The Linux toggle is no longer a separate file. `scripts/toggle-linux.sh` was
merged into the repo-root `toggle.sh`, which dispatches on `$OSTYPE`: an early
`if [[ "$OSTYPE" != darwin* ]]` block runs the self-contained Linux toggle
(shuf / ss / nmcli / ip / systemd) and exits before the macOS body. The shared
DNS/proxy/daemon primitives it used were already moved to `common.sh` in the
`313dc32` unification, so both branches call them by the same names.

---

### recover.sh  Linux path (unified into the root cross-platform script)

Linux recovery is no longer a separate file. `scripts/recover-linux.sh` was
merged into the repo-root `recover.sh`, which dispatches on `$OSTYPE`: the Linux
branch runs a self-contained teardown (resolvectl / nmcli / ip) and exits before
the macOS dual-mode body. All formatting/timeout helpers come from `common.sh`.

---

### scripts/setup-linux.sh (416 lines, 18KB)

**Functions (4 total):**

| # | Function | Line | Purpose |
|---|----------|------|---------|
| 1 | `step()` | L10 | Print formatted step header |
| 2 | `ok()` | L11 | Print green success message |
| 3 | `warn()` | L12 | Print yellow warning message |
| 4 | `die()` | L13 | Print red error + exit 1 |
```

```
**~95% identical to macOS setup.sh.** Only differences:
1. Desktop shortcuts (L366-391 Linux `.desktop` files vs macOS `.app` via `osacompile`)
2. Final instructions text
3. .env template: macOS adds `GATEWAY_USE_EXIT_NODE`, `WARP_FAIL_THRESHOLD`, `KILL_SWITCH`  Linux omits these

---

## Part 3: Utility Scripts

### scripts/diagnose-host-leak.sh (722 lines, 38KB)

**Purpose:** Comprehensive host-routing diagnostic. 8 checks classify host IP leaks. Supports `--fix` and `--watch` modes.

**Duplicated logic:**
- `resolve_active_service()`  identical pattern as toggle.sh, recover.sh, fix-docker-bridge-collision.sh
- `run_with_timeout()`  reimplemented timeout
- Color codes  RED/GREEN/YELLOW/BOLD/NC with tty detection
- WARP check via `cloudflare.com/cdn-cgi/trace`
- `export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"`
- .gateway_ip reading pattern

### scripts/fix-docker-bridge-collision.sh (402 lines, 18KB)

**Purpose:** Fixes Docker bridge subnet collisions with `10.200.1.0/24`.

**Duplicated logic:**
- `resolve_active_iface()`  same interface detection as diagnose-host-leak.sh
- Color codes, step()/ok()/warn()/fail()/die() formatting helpers

### scripts/device-scramble.sh (159 lines, 7.8KB)

**Purpose:** Randomizes hostname and MAC address for Wi-Fi anonymity. Includes fail-safe MAC test and disassociate-first logic.

**Duplicated logic:**
- `need_root()`  same sudo check as other scripts
- `WARP_INHIBIT` marker  interacts with watcher.sh to prevent gateway shutdown during Wi-Fi blip

### scripts/fix-ssh-delay.sh (63 lines, 2.7KB)

**Purpose:** One-shot fix for ~20s SSH delay. Standalone, minimal duplication (uses hardcoded / instead of color functions).

### scripts/reinstall-host-tailscale.sh (740 lines, 31KB)

**Purpose:** Complete uninstall + reinstall of host Tailscale.

**Duplicated logic:**
- `with_timeout()`  yet another bash timeout implementation (polling-based, different from diagnose and toggle versions)
- Color codes + logging functions
- System Extension detection logic (similar to diagnose-host-leak.sh)

### scripts/routing-fix.sh (266 lines, 10.5KB)

**Purpose:** Runs INSIDE warp container. Maintains routing table 200, FORWARD rules, country blocking, IP blocklists, tunnel health checks.

**Duplicated logic:**
- (Fixed) WARP_ENDPOINT defaults are now centralized in common.sh
- WARP health check via cdn-cgi/trace (another instance)
- iptables dual-backend pattern (for-loop over `iptables iptables-legacy`)

### scripts/sweep.sh (491 lines, 28KB)

**Purpose:** Read-only post-restart health sweep  answers "after a toggle/restart/wake, is the gateway actually healthy end-to-end  including under real load?" Never mutates routing, PF, DNS, or containers, so it is safe to run against a live gateway at any time (including from a remote session). `errexit` is deliberately OFF: probes are *expected* to fail on an unhealthy gateway and each exit status is recorded as the diagnostic signal.

**The 7 checks** (each recorded PASS/WARN/FAIL; exit 0 only when nothing FAILs, so it can gate CI/launchd/post-restart hooks):
1. **Containers**  `docker compose ps`: warp/adguardhome/tailscale/tor/routing-fix up (warp + adguardhome additionally `healthy`).
2. **Double-tunnel/leak**  container WARP egress IP == host egress IP, both `warp=on` (gluetun ships wget not curl, so the in-container trace falls back to wget).
3. **Hostcontainer path**  the gateway peer line in `tailscale status` shows `direct <ip:port>`; WARNs on a DERP relay.
```

4. **AdGuard filtering** `compiled_rules.txt` block counts present + live interception (`dig @` the gateway IP from `.gateway_ip`).
5. **PF kill-switch** `pfctl` reports Enabled AND anchor `com.apple/nullexit` carries rules (needs passwordless sudo, else WARN).
6. **Tailnet reachability** (added 2026-07-17; upgraded to burst-ping 2026-07-18) every peer `tailscale status` reports online (self and `offline` peers skipped; a `-` status means online-but-idle and is kept) is sent a *burst* of `tailscale ping --timeout=3s --c $SWEEP_PING_BURST --until-direct=false` (default 8; TSMP probes the WireGuard data path even on devices that drop ICMP, answered by the Tailscale client itself with no app/listener needed, so it works identically over DERP or direct P2P). A single ping only proves "reachable right now"; the burst additionally reports loss % and RTT min/avg/max/jitter, which is what actually predicts whether a call/video/QUIC session on that device holds up a peer can be 100% reachable and still have real quality problems a single ping would hide (confirmed live: an S24 over DERP showed 0% loss but 665ms of RTT jitter). 0% loss PASS; some loss (<40%) WARN; 40% loss or no pong FAIL. DERP-relayed peers are also flagged inline as throughput-degraded, pointing at check 7.
7. **Throughput** (added 2026-07-18) pulls a real file over HTTPS from a fast CDN (Hetzner; Cloudflare's own speed-test endpoint 403s WARP's shared/CGNAT egress IPs, so it's avoided) twice: once over the host's normal double-tunnel route, once through the `warp` container's own SOCKS5 proxy (port read from the live `/tmp/nullexit-ports.env`) to isolate WARP from the Tailscale-wrap hop without exec'ing into or installing anything in the container. Verdict is read from curl's actual exit code, not a byte-count guess: `56` (`Recv failure: Connection reset by peer`) is a real mid-transfer stall (WARN see devref.md 15.3.1); `28` is just the bounded test window ending on an otherwise-healthy transfer (PASS). A slow but supposedly "fast" CDN target was deliberately avoided after `speedtest.tele2.net` (the initial choice) turned out to be the bottleneck itself, not nullexit see devref.md 15.3.6.

Output: coloured stdout summary + a timestamped plain-text report `sweep-<TIMESTAMP>.txt` in the repo root (one file per run, so runs can be diffed); `--quiet` prints only the verdict. Signed in `.signatures` re-run `bash scripts/crypto.sh --sign` after editing it.

scripts/watcher.sh (320 lines, 15KB)

Purpose: Long-running launchd daemon (`com.nullexit.wake-recovery`) for post-wake/post-roam recovery. Runs independently of the gateway. Four backgrounded listeners: (1) WAKE via `log stream`, (2) NETWORK via `scutil n.watch`, (3) `TUNNEL_FAILED_CLOSED.marker` poll, (4) **Wi-Fi uplink keepalive** (added 2026-07-23).

Key function `detect_lan_p2p_mode()`:
Added July 10, 2026. Called on startup and on every network state change (Listener 2 NET: events). Determines whether the current Wi-Fi network safely allows direct Tailscale LAN P2P connections:
1. **WPA2-Enterprise check:** `system_profiler SPAirPortDataType` reads the `Security:` field for the current association. Enterprise = 802.1x = always AP-isolated sets `allow=false`.
2. **AP isolation probe:** `route -n get 1.1.1.1` finds default gateway IP, then `ping -c 3 -t 1 -q "$gw"` if 0 responses on a non-empty network, AP isolation is active sets `allow=false`.
3. **Output:** Writes `'true'` or `'false'` to `.lan_p2p_detected` in the repo root.
4. **Consumer:** `routing-fix.sh` reads this file every 30 seconds and enforces/relaxes the RFC1918 DROP rule accordingly no restart required.

Listener 4 Wi-Fi uplink keepalive (added 2026-07-23, devref 15.5.9):
Polls the Wi-Fi interface every 15s. When it has no routable IP (empty or `169.254.*`) `bash scripts/wifi-rejoin.sh rejoin`; when it has one `... reset`. This is the always-on driver for Wi-Fi auto-rejoin the WARP Watcher (`toggle.sh`) that used to be the only driver runs solely while the gateway is up, so a Wi-Fi drop with the gateway down/off left the Mac stranded. Leak-safe: only acts when there is no routable address (no egress path).

Duplicated logic:
- PATH export (similar to toggle.sh/recover.sh)
- Marker file pattern (`/tmp/nullexit-gateway-active.marker`)

scripts/wifi-rejoin.sh (130 lines) unsigned helper

Purpose: Re-associates ONLY user-remembered Wi-Fi networks (never arbitrary APs). Subcommands: `set`/`forget`/`list` manage `.last_wifi_ssid` (gitignored, one SSID per line); `rejoin`/`reset` are internal, called by `watcher.sh` Listener 4 (and the WARP Watcher while the gateway is up). `rejoin` enforces a 15s grace, progressive backoff (30cap 300s), and a give-up cap (`MAX_TRIES=6`) via `/tmp/nullexit-wifi-*` markers, and joins by name via `sudo -n networksetup -setairportnetwork` (Keychain password). No Location, no live-SSID read (devref 15.11.10). NOT in the crypto `.signatures` list.

scripts/dns-proxy.py (92 lines, 2.9KB) Clean, standalone

scripts/logger.py (144 lines, 5.9KB) Clean, standalone

Part 4: Cross-Cutting Issues

Functions Identical Across macOS Linux

Function	macOS toggle	macOS recover	macOS setup	Linux toggle	Linux recover	Linux setup
`run_with_timeout()`				identical	identical	
`stop_local_dns_proxy()`				identical		
`start_local_dns_proxy()`				~95% similar		
`is_gateway_active()`				~80% similar		

```
| `step/ok/warn/die` |  |  |  |  |  identical |  identical |

### Platform-Adapted Functions (same logic, different OS commands)

| Function | Linux Command | macOS Command |
|---|---|---|
| `get_active_service()` | `ip route get 1.1.1.1 \| awk '{print $5}'` | `route get default` + `networksetup -listnetworkserviceorder` |
| `reset_dns()` | `resolvectl revert` | `networksetup -setdnsservers` |
| `cleanup_network_state()` | `ip link set dev down/up` | `networksetup -setairportpower off/on` + proxy disable |
| `force_dns_to_gateway()` | `resolvectl dns` + `resolvectl domain ~ts.net` | `networksetup -setdnsservers` + `-setsearchdomains ts.net` |
| `start_sleep_prevention()` | `systemd-inhibit --what=sleep` | `caffeinate -i` |
| `enable/disable_socks_proxy()` | Stub (no-op on Linux) | `networksetup -setsocksfirewallproxy` |
| DNS cache flush | `resolvectl flush-caches` | `dscacheutil -flushcache` + `killall -HUP mDNSResponder` |

### Features Missing From Linux Scripts

These exist in `toggle.sh`'s macOS branch but NOT its Linux branch:

1. `write_gateway_active_marker()` / `clear_gateway_active_marker()`
2. `restart_tailscaled_daemon()`
3. `start_warp_watcher()` / `stop_warp_watcher()`  WARP liveness monitor
5. `add_warp_bypass_routes()` / `remove_warp_bypass_routes()` / `setup_exit_node_routing()`
6. `LOG_FILE` structured logging (timestamps to output.log)
7. MSS clamping injection (step 4c)
8. Rule count reporting
9. `--post-wake` mode in recover
10. Container teardown on failure in cleanup_handler

---



## Part 5: Constant Duplication Map

| Value | Files |
|---|---|
| `162.159.192.1` (WARP endpoint) | docker-compose.yml, routing-fix.sh (now dynamically loaded via .env/common.sh!) |
| `5335` (AdGuard DNS port) | docker-compose.yml, AdGuardHome.yaml, post-rules.txt |
| `5354` (mapped DNS port) | docker-compose.yml, dns-proxy.py |
| `41642` (Tailscale WireGuard) | docker-compose.yml, routing-fix.sh, post-rules.txt |
| `1080` (SOCKS5 port) | docker-compose.yml, toggle.sh |
| `1120` (MSS clamp) | post-rules.txt, .env |
| `admin:nullexit` (AdGuard creds) | AdGuardHome.yaml (bcrypt) |
| `10.200.1.0/24` (Docker subnet) | (Fixed) no longer duplicated |
| `America/New_York` (timezone) | docker-compose.yml (4 services) |
| `/tmp/nullexit-caffeinate.pid` | (Fixed) centralized in common.sh |
| `/tmp/nullexit-dns-watcher.pid` | (Fixed) centralized in common.sh |
| `/opt/homebrew/bin:...` (PATH) | (Fixed) centralized in common.sh |

---

## Part 6: Refactoring Outcome (Implemented July 2026)

**Decision: "Scripts-only" architecture**
Instead of introducing a new `lib/` directory, all shared code was moved into the existing `scripts/` directory to keep the project hierarchy flat and simple.

### 1. `scripts/common.sh` (288 lines, 10KB)
Extracts the fundamental primitives and networking logic used across orchestrator scripts:
- **Formatters:** `step()`, `ok()`, `warn()`, `fail()`, `die()`
- **Timeout Wrapper:** `run_with_timeout()` (Canonical version with PID tracking, graceful SIGTERM, and SIGKILL fallback)
- **Daemon Lifecycle:** `restart_tailscaled_daemon()`, `stop_pidfile_daemon()` (collapsed caffeinate, DNS watcher, and WARP watcher teardown logic)
- **Network Interface/Routing:** `get_active_service()`, `get_en0_service()`, `bounce_wifi_interfaces()`, `add_warp_bypass_routes()`, `remove_warp_bypass_routes()`
- **Proxy/DNS Cleanup:** `disable_all_proxies()`, `flush_dns_cache()`, `reset_sharing_services()`
- **Configuration Helpers:** `read_adguard_ip()`, `is_kill_switch_enabled()`, `get_warp_endpoint_1()`, `get_warp_endpoint_2()` (which centralizes the hardcoded 162.159.192.1 / .193.1 into .env logic)

### 2. `scripts/setup-common.sh` (~350 lines)
Extracts the heavy, duplicated installation logic shared between `setup.sh` (macOS) and `scripts/setup-linux.sh`:
- Docker and Colima prerequisite checks
- `wgcf` binary download and WARP key generation
- `.env` configuration file generation
- Interactive Tailscale Auth Key and AdGuard Rule Profile prompts
```

```
### Sourcing Pattern
- macOS root scripts (`toggle.sh`, `recover.sh`, `setup.sh`) source via:
  `source "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/common.sh"`
- Scripts under `scripts/` (`watcher.sh`, `diagnose-host-leak.sh`, etc.) source via:
  `source "$(dirname "${BASH_SOURCE[0]}")/common.sh"`

## Agent Verbs / Protocols

### `/sweep` (or `sweep the gateway`)
When the user invokes this verb, the agent MUST perform a full end-to-end connectivity, health, security, and performance audit of the gateway:
1. **WARP Validation:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp=` (must equal `on`).
2. **Tor Proxy & Control Validation:**
   - Source the dynamically generated ports from `/tmp/nullexit-ports.env` (or identify them using `docker compose ps`).
   - **SOCKS5 Check:** Run `curl -s --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip` to confirm the SOCKS5 proxy successfully connects and routes through the Tor network.
   - **Control Port Check:** Query the Tor Control Port using netcat: `echo "PROTOCOLINFO" | nc 127.0.0.1:$TOR_CONTROL_PORT`. Verify it returns a standard `250-PROTOCOLINFO` response, confirming the Control Port is active and communicating.
   - **Circuit & Node Lookup:** Query the Tor Control Port for the active path (`GETINFO circuit-status`). For each active built circuit, parse the relay fingerprints, lookup their IP addresses (`GETINFO ns/id/<fingerprint>`), and resolve their country codes using Tor's local GeoIP database (`GETINFO ip-to-country/<IP>`). Include the list of active circuits, showing the role (Guard/Middle/Exit), nickname, IP, and country of each hop in the final sweep report.
   - **Transparent Onion Validation:** Run `nslookup duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion` to verify it resolves to an IP within `198.18.0.0/15` benchmark range. Verify transparent dark web routing by running `curl -sI -H "Host: duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion" http://<resolved-ip>/` (should return HTTP 301/200).
3. **DNS Leak Validation:**
   - Audit the upstream resolver currently utilized by the host. Run `curl -s https://edns.ip-api.com/json` to fetch details of the DNS resolver processing the client's requests.
   - Verify that the resolver IP and organization returned belong to Cloudflare (e.g., AS13335) or match the WARP tunnel's exit range, and that **no** DNS traffic leaks to the user's raw physical ISP DNS.
4. **Performance Benchmark Audit:**
   - Run the python benchmark script `python3 benchmarks/test_load.py`.
   - Report the overall load speed and the ratio of successfully fetched vs. blocked/failed subresources for the targeted domains to ensure ad-blocking and MTU encapsulation are performing correctly.
5. **Tailscale Mesh Validation:**
   - Run `ping -c 1 100.100.21.8` to ensure the gateway is reachable.
   - Run `tailscale status` to identify all active/online nodes in the mesh.
   - For every online peer returned in the status output, execute `tailscale ping --timeout=3s --c 8 --until-direct=false <peer-ip>` (a burst, not a single ping  TSMP works even on phones that drop ICMP) to verify connectivity and read loss % + RTT spread across the burst, not just one-shot latency. `bash scripts/sweep.sh` automates this as its check 6/7 and reports each peer's path (direct endpoint vs `DERP(...)`), loss %, and RTT min/avg/max/jitter; its check 7/7 additionally measures real throughput (host path vs. a WARP-only baseline).
6. **Log Audit:** Run `tail -n 100 output.log | grep -iE "(error|fail|warn)"` to catch any underlying component failures or warnings.

### `/latex` (or `generate latex`)
When the user invokes this verb, the agent MUST run `python3 -I scripts/generate_tex.py` to regenerate the documentation source code (`nullexit_unified.tex`). The agent MUST NOT attempt to compile the `.tex` file into a PDF (the user handles PDF compilation locally).
```

4 devref.md

```
# nullexit  Development Reference & Resolved Issues

> **Last updated:** July 19, 2026
> **Purpose:** Provide any LLM or developer with complete project understanding, debugging history, and resolved issues so they can make informed changes without re-reading every file.
> **Diagrams:** See [`diagrams.md`](./diagrams.md) for system architecture, toggle.sh flowcharts, monitoring layer, traffic sequence, recover.sh decision tree, and the full failureself-healing map.

This reference is organized into four parts:
> - **Part I  Architecture & Reference** (18)
> - **Part II  Design Decisions & Threat Model** (914)
> - **Part III  Incident Log & Resolved Issues** (15)
> - **Part IV  Observations, Changelog & TODO** (1618)

---

# Part I  Architecture & Reference

```

```
## 1. What Is nullexit?

nullexit is a **Tunnel-in-Tunnel VPN gateway** that chains **Tailscale** (mesh VPN) through **Cloudflare WARP** (exit VPN). It runs on a macOS host inside Docker containers managed by **Colima** (lightweight Linux VM). The goal: any device on the user's Tailscale mesh can use this gateway as an exit node, and all traffic exits through Cloudflare WARP  achieving double encryption, ISP invisibility, and network-wide ad-blocking via **AdGuard Home**.

### Traffic Flow
```
Phone/Laptop Tailscale mesh (WireGuard) Mac host Docker container
   Tailscale decrypts Gluetun re-encrypts WARP tun0 Cloudflare Internet
```

### SOCKS5 Failover Path (if exit node is unavailable)
```
DNS: Browser host 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP DNS
TCP: Apps macOS SOCKS5 proxy (127.0.0.1:1080) container tun0 WARP Internet
```

---

## 2. Architecture

### Containers (all share `warp`'s network namespace via `network_mode: service:warp`)
| Container | Image | Role |
|-----------|-------|------|
| `warp` | `qmcgaw/gluetun:v3.41.1` | Gluetun WireGuard client  Cloudflare WARP. Owns the network namespace. Strict firewall. |
| `tailscale` | `tailscale/tailscale:v1.98.4` | Advertises as exit node on the mesh. Also provides the built-in SOCKS5 proxy fallback via `TS_SOCKS5_SERVER=:1080`. |
| `routing-fix` | `alpine:3.20` | Sidecar that maintains routing tables + iptables rules every 30 seconds. |
| `rule-compiler`| `golang:1.22-alpine` / `alpine:3.20` | One-shot startup container that compiles 500k+ DNS block rules in <2s and immediately exits. |
| `adguardhome` | `adguard/adguardhome:v0.107.77` | DNS sinkhole for ads/trackers. Listens on port 5335. Upstream DNS: `127.0.0.1:53` (through WARP). |

### Port Mappings (on host via Colima)
| Host Port | Container Port | Protocol | Purpose |
|-----------|---------------|----------|---------|
| 5354 | 5335 | TCP+UDP | AdGuard DNS |
| (Tailscale IP):3000 | 3000 | TCP | AdGuard web UI (Internal to Mesh) |
| 41641 | 41641 | UDP | Tailscale WireGuard direct |
| 1080 | 1080 | TCP | SOCKS5 proxy |

### Host Components
- **Colima VM**  Linux VM running Docker (600MB RAM + 400MB swap)
- **tailscaled**  Standalone Tailscale daemon via `brew install tailscale` + `brew services start tailscale` (NOT the GUI `.app`)
- **macOS network settings**  DNS, SOCKS proxy, routing all manipulated by `toggle.sh`

---

## 3. Key Files

| File | Lines | Purpose |
|------|-------|---------|
| `toggle.sh` | ~2109 | **Main script (macOS + Linux).** Dispatches on `$OSTYPE`  an early Linux block (shuf/ss/nmcli/ip, systemd) runs and exits before the macOS body (jot/netstat/networksetup, launchd/pf). Detects state, toggles gateway ON/OFF. Handles DNS hijacking, Tailscale exit node, SOCKS proxy, sleep prevention, timeouts, cleanup. |
| `setup.sh` | 406 | **One-time setup.** Installs deps (Docker, Tailscale, wgcf), generates WARP keys, writes `.env`, configures AdGuard via API, starts containers. |
| `docker-compose.yml` | 194 | Service definitions for all 5 containers. |
| `scripts/routing-fix.sh` | 201 | Maintains routing tables (table 200 for SOCKS5, table 52 for CGNAT) and FORWARD iptables rules in both nftables and legacy backends. Loads and enforces the IP blocklist via `ipset`. Runs in a 30-second loop. |
| `post-rules.txt` | 18 | Gluetun iptables rules loaded at container start. FORWARD accept for tailscale0, NAT MASQUERADE on tun0, DNS redirect to port 5335, IPv6 drop, TCP MSS clamping. |
| `scripts/dns-proxy.py` | 91 | Local DNS proxy. UDP:53  TCP:5354 with proper 2-byte length prefix. Fallback when exit node is unavailable. |
| `scripts/rule-compiler/` | ~800 | Unified threat compiler (ported to Go from Python). **DNS pipeline:** fetches remote blocklists, deduplicates against AdGuard native filters, optimizes subdomains (~50% reduction), compiles to AdGuard syntax. **IP pipeline:** fetches threat-intel feeds (Spamhaus, Feodo, ET, CINS), strips private/Tailscale ranges, outputs atomic `ipset` restore file. Built dynamically in Docker multi-stage build. |
| **`adguard/work/`** |  | **Bind-mounted container data folder. Off-limits from outside Docker  READ-ONLY-EXCEPT-VIA-`sync-rules`. See README 6.** |
| `recover.sh` | ~742 | Nuclear recovery script (**macOS + Linux**  dispatches on `$OSTYPE`; the Linux branch runs a self-contained teardown via `resolvectl`/`nmcli`/`ip` and exits, macOS falls through to the dual-
```

```
    mode body). Gain: use `--post-wake` (macOS) for non-destructive refresh after sleep/wake or Wi-Fi roam
    keeps the gateway live while still re-hijacking DNS + refreshing Tailscale exit-node + force-recreating
    warp if unhealthy. Resets DNS on ALL services, disables proxies, flushes routes, stops containers, kills
    sleep prevention, power-cycles Wi-Fi. Cryptographically verified. |
| `scripts/diagnose-host-leak.sh` | 580+ | **One-shot host-routing diagnostic.** Runs 8 checks (Tailscale,
    SOCKS5, CDN-cgi/trace, IPv6 leak, default route, output.log hints, gateway IP, Tailscale System Extension
    conflict), classifies into ONE of three known leak scenarios (A=SOCKS5 fallback, B=IPv6 leak, C=route
    freeze) or OK, writes a timestamped report to `host-leak-diagnostic-<UTC>.txt`, and prints a ready-to-run
    fix on stdout. Pass `--fix` to apply the matched remediation and re-verify egress. Pass `--watch` (or `--
    watch 30`) to run a full baseline then continuously monitor warp/IPv6/default-route every N seconds,
    alerting on any state change. See 15.6.1. |
| `scripts/host-leak-probe.sh` |  | **REMOVED (2026-07-15, verified deleted/untracked).** Was a continuous sub-
    second host-egress prober auto-launched via `HOST_LEAK_PROBE`. Its fail-closed role is now covered by the
    **PF kill-switch** (`enable_killswitch` in `common.sh`) + the in-container **WARP Watcher** (15.6.2); on-
    demand auditing is `scripts/diagnose-host-leak.sh`. See 15.6.1 (retained as historical rationale). |
| `scripts/fix-docker-bridge-collision.sh` | ~290 | **One-shot fix for 15.2.3 (Docker bridge subnet collision)
    .** Detects Docker's `172.17.0.0/16` bridge colliding with the host Wi-Fi subnet (the symptom that produces
     `default 172.17.0.1 UGScg en0` in `netstat -rn`), picks a non-colliding `docker.bip` from a small
    candidate set, idempotently edits `~/.colima/default/colima.yaml` with a timestamped backup, restarts
    Colima, rebinds Wi-Fi DHCP, verifies the new default route is no longer Docker's bridge, and re-runs `
    toggle.sh` + `scripts/diagnose-host-leak.sh`. Supports `--dry` and `--skip-toggle`. |
| `scripts/watcher.sh` | 320 | **Long-running post-wake + post-roam daemon.** Launched by launchd LaunchAgent;
    subscribes to `com.apple.powermanagement` events via `log stream` and to network-state changes via `scutil
    n.watch`. Both fire `recover.sh --post-wake`. A 4th listener (added 2026-07-23) polls the Wi-Fi interface
    and drives `scripts/wifi-rejoin.sh` when the uplink has no routable IP  independent of gateway state
    (15.5.9). Single-instance lock + SCRIPT_DIR-relative resolve of `RECOVER`. See 15.5.1. |
| `scripts/wifi-rejoin.sh` | 130 | **Wi-Fi keepalive helper (unsigned).** Re-associates ONLY networks the user
    remembered in `.last_wifi_ssid` (`set`/`forget`/`list`), never arbitrary APs. `rejoin` enforces a grace
    period, progressive backoff, and a give-up cap; `reset` clears the down-timer once the link returns. No
    Location, no root plist read, no live-SSID read (15.11.10). Driven by `watcher.sh` (always-on) and, when
    the gateway is up, the WARP Watcher; the two share `/tmp` markers. See 15.5.9. |
| `scripts/common.sh` | ~40 | **Shared script primitives.** Centralizes formatted echo statements (`step`, `ok
    `, `warn`, `die`) and the safe background `run_with_timeout` execution wrapper. Sourced by all orchestrator
     scripts across macOS and Linux. |
| `scripts/setup-common.sh` | ~350 | **Shared installation logic.** Centralizes logic for verifying Docker,
    installing Tailscale/wgcf, generating `.env`, and prompting for auth keys. Sourced by `setup.sh` and `setup
    -linux.sh`. |
| `scripts/setup-linux.sh` | ~80 | **Linux sibling of `setup.sh`.** Sources `setup-common.sh`. Distro detection
     (apt/dnf/pacman), installs Linux-specific dependencies, creates `.desktop` shortcuts. |
| `scripts/Toggle-Gateway.bat` | ~300 | **Windows Toggle helper.** Moved from root to `scripts/`. Helps invoke
    the framework on Windows if applicable. |
| `scripts/reinstall-host-tailscale.sh` | ~740 | **Nuke-and-reinstall tool for host Tailscale.** Completely
    purges Tailscale, its settings, macOS Background Items (`.btm` database), and stale System Extensions.
    Wipes ALL Login Items via `sfltool reset-login-items` (requires sudo). Useful for crisis recovery but
    destructive to other login items. |
| `scripts/fix-ssh-delay.sh` | ~100 | **Fixes SSH connection delays.** One-shot script to address SSH delays
    caused by DNS or routing issues, useful in crisis situations. |
| `scripts/logger/` | ~100 | **Internal logger piped from `scripts/routing-fix.sh`** inside the `warp`
    container. Ported to Go from Python. Built via Docker multi-stage build. |
| `black_list.txt` | 86 | Custom domains to block (ads, trackers, telemetry, cross-platform device-
    identification endpoints). Supports `$important` modifier. |
| `white_list.txt` | 231 | Domains to force-allow (YouTube, Apple services, etc.). Always wins over blocks. |
| `.env` | ~14 | WARP WireGuard keys, Tailscale auth key, rule profile. **Contains secrets & NULLEXIT_SEED.** |
| `.gateway_ip` | 1 | Static gateway Tailscale IP (fallback for dynamic resolution). |
| `scripts/unlock-files.sh` | ~20 | **One-shot stale-permission fix.** Uses atomic rename (`cp` + `mv`) to
    replace locked inodes (mode `000`/`0444` from old `chmod` decisions) with fresh writable ones  no `chmod`
    called. |
| `scripts/crypto.sh` | ~50 | **Cryptographic integrity enforcement.** Uses `NULLEXIT_SEED` from `.env` to sign
     core bash scripts (HMAC-SHA256) and strictly verify them on start to prevent manipulation. Pass `--sign`
    or `--verify`. |
| `scripts/pf.conf` | 35 | **macOS native firewall ruleset.** Enforces the default-deny kill-switch, allows
    local LAN / Tailscale traffic, and applies TCP MSS Clamping (`max-mss 1160`) to prevent WireGuard MTU
    fragmentation. |

---

## 4. toggle.sh Flow

### State Detection
`is_gateway_active()` returns true if EITHER containers are running OR host DNS is not `1.1.1.1`. This dual
    check prevents the script from starting when it should stop (e.g., containers crashed but DNS is still
    hijacked).

### START Path (gateway OFF  ON)
1. **Reset DNS to 1.1.1.1**  Prevents deadlocks during startup
2. **Disconnect host Tailscale**  `tailscale down` (prevents exit-node routing during container boot)
3. **Compile DNS rules**  `docker compose run --rm rule-compiler` (Go binary inside multi-stage Docker build)
4. **Boot Colima VM** (if not running)  via `colima_start_until_ready()`, which returns as soon as Docker is
    reachable (~12s) instead of blocking on Colima's ~2-min post-provision hang (15.8.7). The networking mode
    is **gated on the LAN-P2P signal** (`.lan_p2p_detected`, falling back to `TAILSCALE_ALLOW_LAN_P2P` in `.env
    `): a trusted home network requests `--network-address --network-mode bridged` (or `shared` on WPA2-
```

```
     Enterprise), while the common AP-isolated / untrusted case boots plain fast user-mode networking (`colima
     start --memory 0.6 --vm-type vz`).
107    - **Why Bridged Networking (when enabled)?** The `--network-mode bridged` and `--network-address` flags
       force the VM to receive its own IP directly from your local router's DHCP server, placing it on the same
       physical subnet as the host. This bypasses macOS network isolation and is what lets **external LAN devices
       ** (phones, laptops on the same LAN), mDNS, and local service discovery reach across the container boundary
       . It requires the `socket_vmnet` helper and is only worthwhile on a trusted LAN, which is why it is gated
       rather than always requested. **It is *not* required for direct hostcontainer P2P**  that path (the gateway
         its own Mac) works even in user-mode networking via the routing-fix `.host_ips` RETURN-rule carve-out;
       see 15.3.5 and 15.8.7.
108    - **Enterprise Wi-Fi Fallback & Dynamic Roaming:** `toggle.sh` dynamically queries `system_profiler
       SPAirPortDataType` on boot. If it detects a **WPA2-Enterprise** connection (802.1X), it forces Colima to
       fall back to `--network-mode shared`. This prevents aggressive network intrusion systems on corporate/
       university Wi-Fi switches from terminating the host's Wi-Fi connection due to "MAC spoofing" (detecting
       multiple MAC addresses originating from a single authenticated port). P2P connections are impossible on
       these networks anyway due to AP Client Isolation.
109      - **Dynamic Roam Downgrading:** `toggle.sh` writes its selected network mode to `/tmp/nullexit_colima_mode
       .txt`. When roaming between networks (e.g., Home to University), `recover.sh` actively checks this state
       against the new network's security requirement (`.lan_p2p_detected`). If a mismatch is detected (e.g.,
       bridged Colima on an Enterprise network), `recover.sh` automatically forces a full `toggle.sh --restart` in
        the background to safely transition the Virtual Machine's network mode and protect the host from de-
       authentication.
110 5. **Configure VM swap**  400MB swap file inside the VM to prevent OOM
111 6. **Clean corrupted AdGuard config**  Remove empty `AdGuardHome.yaml`
112 7. **Start containers**  `docker compose up -d`
113 8. **Wait for gateway Tailscale**  Poll `tailscale status` for "offers exit node" (up to 60s, abort at 40
     consecutive NoState)
114 9. **Resolve gateway IP**  From `.gateway_ip` or `docker compose exec tailscale tailscale ip -4`
115 10. **Connect host to mesh**  Verify `tailscaled` is running (auto-start if needed), then `tailscale up --reset
      --ssh=true --accept-dns=false --accept-routes=true --exit-node=` (`--accept-routes=true` must be explicit
      `--reset` reverts it to `false` otherwise, silently preventing the default route from switching to `utun
     *`)
116 11. **Pre-flight checks**  [1/3] `tailscale ping` gateway, [2/3] `dig +tcp` AdGuard DNS, [3/3] WARP container
     internet
117 12. **If all pass**  Set exit node + hijack DNS to gateway IP (single server, no fallback)
118 13. **If any fail**  Enable SOCKS5 proxy + local DNS proxy as fallback
119 14. **Final DNS re-force**  `force_dns_to_gateway` called again after exit-node transition (tailscaled clobbers
      DNS briefly)
120 15. **Enable sleep prevention**  `caffeinate -i` in background to prevent idle sleep
121
122 ### STOP Path (gateway ON  OFF)
123 1. Disconnect host Tailscale (`tailscale down`)
124 2. `docker compose down -t 5`
125 3. Reset DNS to 1.1.1.1
126 4. Stop local DNS proxy
127 5. **Stop sleep prevention**  Kill caffeinate process
128 6. Full network state cleanup (proxies off, DNS cache flush, route flush, Wi-Fi power cycle)
129
130 ### Safety Mechanisms
131 - `run_with_timeout()`  Pure-bash watchdog for every system command (15s120s)
132 - `cleanup_handler()`  Trap on ERR/INT/TERM/HUP restores DNS to 1.1.1.1, kills caffeinate, and tears down
     Tailscale
133 - `force_dns_to_gateway()`  Sets DNS and VERIFIES via `networksetup -getdnsservers` (3 attempts with backoff)
134 - `SUCCESS_RUN` flag  Cleanup only runs if script didn't complete successfully
135 - `start_sleep_prevention()` / `stop_sleep_prevention()`  Manages `caffeinate -i` via PID file at `/tmp/
     nullexit-caffeinate.pid`
136
137 ---
138
139 ## 5. Routing & Firewall Architecture (Inside Container)
140
141 ### IP Rules (`ip rule show`)
142 | Priority | Match | Table | Purpose |
143 |----------|-------|-------|---------|
144 | 99 | `to 100.64.0.0/10` | 52 | CGNAT range  Tailscale routes (injected by routing-fix) |
145 | 100 | `from 172.18.0.2` | 200 | Container's own traffic (SOCKS5 proxy)  tun0 |
146 | 101 | all | 51820 | Gluetun's WireGuard fwmark  tun0 |
147
148 ### Table 200 (SOCKS5 proxy traffic)
149 - `162.159.192.1 via $DOCKER_GW dev eth0`  WARP endpoint exception (prevents tunnel loop)
150 - `default dev tun0`  Everything else through WARP
151
152 ### iptables (TWO separate stacks!)
153 - **nftables backend** (`iptables`)  Used by Gluetun's `post-rules.txt`. FORWARD policy DROP.
154 - **legacy backend** (`iptables-legacy`)  Used by Tailscale's `ts-forward`. FORWARD policy ACCEPT.
155 - Both stacks apply to every packet. FORWARD RELATED,ESTABLISHED must exist in BOTH.
156
157 ### post-rules.txt (loaded by Gluetun)
158 ```
159 FORWARD: ACCEPT for tailscale0 in/out
160 INPUT: ACCEPT for tailscale0
```

```
NAT POSTROUTING: MASQUERADE on tun0
MANGLE: TCP MSS clamp to 1180
NAT PREROUTING: Redirect DNS (53)  5335 on tailscale0 and eth0
ip6tables: DROP FORWARD on tailscale0
```

5.1 Tailscale Local Network Discovery (DERP Bypass)
By default, Gluetun's strict NAT swallows all of Tailscale's UDP hole-punching packets, forcing Tailscale to fall back to incredibly slow DERP relay servers (often adding 500ms+ latency).

To fix this, we use `FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8` in `docker-compose.yml` (on the `warp` container). This creates `ip rule` bypasses (Priority 99, Table 199) that route all packets destined for private IP ranges *outside* the WARP tunnel directly to the host's `eth0`.
- **172.16.0.0/12** is especially critical because many modern Wi-Fi routers (and Docker itself) assign IPs in this block (e.g., `172.17.x.x`).
- This bypass allows Tailscale's UDP packets to reach devices on the same Wi-Fi directly, establishing a fast peer-to-peer connection while the actual web traffic inside the tunnel remains fully routed through WARP.

5.1.1 Why This Matters: Untrusted, Non-Enterprise Networks

Most public/consumer hotspots hotels, cafes, airports, a friend's home router are **untrusted** (any other guest is a stranger, no admin-enforced isolation) but **not** enterprise-hardened the way WPA2-Enterprise/802.1x campus networks are (see 15.7.1's SNAT poisoning analysis for why *those specifically* get P2P disabled). Because open networks usually lack AP client isolation, `watcher.sh`'s `detect_lan_p2p_mode()` probe (a ping to the default gateway see 16.4) typically finds them P2P-capable and flips `TAILSCALE_ALLOW_LAN_P2P` on automatically.

This is the scenario where nullexit's mesh-first design pays off over a conventional single-device VPN app:

1. **Encryption doesn't depend on P2P succeeding.** Every packet between mesh devices is WireGuard-encrypted whether it goes direct or via DERP. An open network's *lack* of isolation is itself the risk anyone else on that hotel Wi-Fi can reach your device at L2 but it doesn't expose your traffic's content, only its existence/metadata to that local segment.
2. **Zero configuration per network.** The device only needs to be Tailscale-enrolled once, ever. `watcher.sh` re-runs its isolation probe on every Wi-Fi change (`scutil n.watch` on `State:/Network/Global/IPv{4,6}`, 15.5.1) and rewrites `.lan_p2p_detected` accordingly, which `routing-fix.sh` picks up within 30s no app to reopen, no profile to import, no manual toggle per hotspot. The same phone gets DERP-only (safe, no SNAT risk) on an isolated campus network and fast direct P2P (safe, isolation genuinely absent) at a hotel that night, entirely automatically.
3. **The auto-detection is what makes this safe to leave alone.** A user who manually forced `TAILSCALE_ALLOW_LAN_P2P=true` globally and never touched it again would carry the P2P attempt onto the *next* enterprise network they visit, reproducing the 15.7.1 SNAT-poisoning blackout. The watcher's per-network re-probe is what lets the setting behave correctly on untrusted-but-open networks without that risk following the user everywhere.

5.2 Firewalling & Per-Device Access Control
Because every mesh device's traffic passes through the `warp` container's network namespace, this namespace serves as the ultimate choke point for network-wide firewall rules.

Where to Put Rules
The right place to add firewall rules is `scripts/routing-fix.sh`. It runs a 30-second loop inside the namespace and already handles re-injecting rules that Gluetun resets on reconnect. **Do not use `post-rules.txt` for dynamic rules**, as Gluetun flushes and resets it upon VPN reconnects.

The Dual iptables Backend Constraint (CRITICAL)
The container runs two simultaneous iptables backends: `iptables` (for the nftables backend used by Gluetun) and `iptables-legacy` (for Tailscale). Every rule **must** be added to both stacks, or it will silently fail for certain traffic.

Avoiding Duplication
Because `scripts/routing-fix.sh` runs in a loop, you must always check for a rule's existence with `-C` before appending with `-A`. Otherwise, your rules will duplicate every 30 seconds.

Per-Device Identification & Filtering
Each mesh device has a stable `100.x.x.x` Tailscale IP, which is visible as the `src` IP in the `FORWARD` chain. This is how you identify devices for filtering.
* **Domain-level:** AdGuard Home (port 3000, credentials `admin/nullexit`) provides a REST API that supports per-client blocking rules based on their Tailscale IP.
* **IP-level:** Use `iptables` in `scripts/routing-fix.sh` (e.g., `iptables -I FORWARD -s 100.x.x.x -d <blocked_ip> -j DROP`).
* **Country-level:** Add `ipset` to the `routing-fix` apk install step in `docker-compose.yml`, and load CIDR ranges from `ipdeny.com` into an ipset, then block that set in the `FORWARD` chain. See 5.3 (Geo-IP Blocking) for the full configuration flow.

5.3 Geo-IP Blocking
To block traffic to and from specific countries, the system dynamically downloads country-specific IP ranges from [ipdeny.com](http://www.ipdeny.com/ipblocks/data/countries/) and drops them via iptables `FORWARD`

```
    rules.

To add a new country to the blocklist:
1. Find the 2-letter ISO country code (e.g., `cn` for China, `ru` for Russia, `il` for Israel, `kp` for North Korea).
2. Open your `.env` file.
3. Update the `BLOCKED_COUNTRIES` variable: e.g., `BLOCKED_COUNTRIES="kp il cn ru"`.
4. Run `toggle.sh` to restart the gateway and apply the new environment variables to the routing-fix container.

> [!NOTE]
> **Limitations of Geo-IP Blocking:** IP-based blocking is highly effective at blocking servers, direct apps, and raw infrastructure physically located and registered in a specific country (e.g., `www.gov.il`). However, it will **not** block high-profile websites (e.g., `jpost.com`) that route their traffic through global Content Delivery Networks (CDNs) like Google Cloud, Cloudflare, or Fastly. Because the CDN's IP addresses are registered in the US or globally, the network traffic physically never routes to the blocked country, allowing it to bypass the Geo-IP firewall.

## 6. macOS Quirks to Remember

1. **Colima SSH tunnel is TCP-only**  UDP ports (like 5354/udp) are NOT accessible from host. Use `dig +tcp` or the Python DNS proxy (UDPTCP converter).
2. **`docker compose exec -T` injects `\r`**  Always `tr -d '\r'` when capturing output for comparisons or `networksetup` commands.
3. **`brew services` can get stuck**  Fix with `sudo brew services restart tailscale`. Common state: `brew services list` shows `started` but `tailscale status` shows "Tailscale is stopped."
4. **No `timeout` command**  macOS lacks GNU `timeout`. Use the `run_with_timeout()` bash function.
5. **`sudo -v` prompts even with NOPASSWD**  Use `sudo -n` (non-interactive) everywhere.
6. **DNS is scoped to en0**  macOS scutil DNS resolver is scoped to en0 (Wi-Fi). DNS changes on other interfaces (like USB Ethernet) are ignored unless en0's service is also updated. The script always sets DNS on BOTH `$ACTIVE_SERVICE` and `$EN0_SERVICE`.
7. **tailscaled clobbers DNS during exit-node transition**  `force_dns_to_gateway` is called a second time at the end of the ENABLE branch to counteract this.
8. **`tailscale up --reset`** resets all unspecified flags to defaults  Every `--reset` call MUST also pass `--accept-dns=false` explicitly or tailscaled re-enables DNS management.
9. **`docker compose ps` CWD pitfall**  `docker compose` commands require a `docker-compose.yml` in the current working directory. Use `docker ps` for raw container listing from any CWD; `docker compose ps` requires the project directory.

---

## 7. Environment Variables

### `.env` File
| Variable | Purpose |
|----------|---------|
| `WIREGUARD_PRIVATE_KEY` | WARP WireGuard private key (from `wgcf generate`) |
| `WIREGUARD_PUBLIC_KEY` | WARP WireGuard public key |
| `WIREGUARD_ADDRESSES` | WARP WireGuard address (e.g., `172.16.0.2/32`) |
| `TS_AUTHKEY` | Tailscale auth key for the container |
| `NULLEXIT_SEED` | 256-bit seed for HMAC-SHA256 script integrity signing (generated by `setup.sh`) |
| `GATEWAY_RULE_PROFILE` | Rule compilation tier: `light`, `medium`, `heavy` |
| `GATEWAY_BYPASS_PING` | (Optional) `true` to proceed even if pre-flight checks fail |
| `GATEWAY_USE_EXIT_NODE` | (Optional) `false` to skip exit node, DNS-only mode |
| `GATEWAY_HIJACK_HOST` | (Optional) `false` to skip DNS hijacking on the host (VPN/adblocking for Tailscale peers only) |
| `GATEWAY_MSS` | (Optional) TCP MSS clamp value. Default and only validated-safe value is `1120`  15.3.2 found 1180 (and 1160) still too large for the double-encapsulation overhead and causes mid-transfer stalls. Applied to both the container (`post-rules.txt`) and the host (`pf.conf`), kept in sync. |
| `WARP_FAIL_THRESHOLD` | (Optional) Consecutive `warp=off` polls before auto-shutdown (default 6 = 30s; 3 = 15s) |
| `WARP_ENDPOINT_1` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 1 (default `162.159.192.1`) |
| `WARP_ENDPOINT_2` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 2 (default `162.159.193.1`) |
| `KILL_SWITCH` | (Optional) `true` to enforce strict PF lock  drops all host traffic if VPN fails. Breaks SSH if VPN dies. |
| `STOP_COLIMA_ON_EXIT` | (Optional) `true` to fully shut down the Colima VM on toggle-off (saves battery on dedicated hosts) |
| `ADGUARD_USER` | (Optional) AdGuard Home web UI username (default: `admin`) |
| `ADGUARD_PASSWORD` | (Optional) Shared password for AdGuard Home and Tor ControlPort (default: `nullexit`) |
| `BLOCKED_COUNTRIES` | Space-separated 2-letter ISO codes to block via ipdeny.com CIDR ranges (e.g. `"kp il cn ru"`) |

---

## 8. Diagnostic Commands

### Container Health
```bash
docker ps --format '{{.Names}} {{.Status}}'
docker compose exec -T tailscale tailscale status
docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace
```
```

```

### SOCKS5 Proxy
```bash
curl --socks5-hostname 127.0.0.1:1080 --max-time 5 -s https://www.cloudflare.com/cdn-cgi/trace
```

AdGuard DNS
```bash
dig +tcp @127.0.0.1 -p 5354 google.com +short +timeout=5
```

Firewall & Routing
```bash
# nftables FORWARD chain
docker compose exec -T warp iptables -L FORWARD -v --line-numbers
# Legacy FORWARD chain
docker compose exec -T warp iptables-legacy -L FORWARD -v --line-numbers
# IP rules + routing tables
docker compose exec -T warp ip rule show
docker compose exec -T warp ip route show table 199
docker compose exec -T warp ip route show table 200
```

Host macOS
```bash
tailscale status
brew services list | grep tailscale
networksetup -getdnsservers Wi-Fi
networksetup -getsocksfirewallproxy Wi-Fi
sysctl net.inet.ip.forwarding
```

Part II Design Decisions & Threat Model

9. Privacy Architecture & Threat Model

Layered Defense
The gateway stacks four layers targeting different attack surfaces:

- **Layer 1 ISP-Level Surveillance (WARP):** Your ISP sees only an encrypted WireGuard tunnel to a Cloudflare IP. In the US, ISPs can legally sell your browsing metadata and are subject to secret NSLs (National Security Letters). WARP removes their visibility entirely.
- **Layer 2 DNS Tracking (AdGuard Home):** DNS is the phone book of the internet. By default, queries go to your ISP's resolver, giving them a complete log of your traffic. AdGuard intercepts all DNS from mesh devices, blocks trackers, and resolves queries through the WARP tunnel.
- **Layer 3 Device Identity & IP Exposure (Tailscale):** On untrusted networks (coffee shops, public Wi-Fi), your device's IP is exposed. Tailscale creates an encrypted WireGuard mesh between your devices, routing all traffic to your gateway exit node.
- **Layer 4 Kernel-Level IP Blocking (ipset/iptables):** Sophisticated malware bypasses DNS entirely with hardcoded IPs. The `rule-compiler` fetches ~16,700 threat-intelligence IPs/CIDRs (Spamhaus, Feodo Tracker, Emerging Threats, CINS) and enforces them as `DROP` rules in the kernel `FORWARD` chain blocking both outbound C2 connections and inbound attack traffic.

The Documented Threat: Why This Matters
This architecture is calibrated against documented mass-surveillance programs revealed in the Snowden disclosures:
- **PRISM:** Bulk collection of communication content from major tech providers.
- **UPSTREAM:** Tapping physical backbone fiber, collecting metadata and content in bulk.
- **XKeyscore:** Retroactive search of unencrypted traffic.
- **MUSCULAR:** Infiltration of unencrypted internal datacenter interconnect links.

Passive bulk collection is not paranoia it is a documented reality. By double-encrypting and routing traffic, nullexit prevents your ISP from harvesting your metadata.

Trust Assumptions & Shifting
Every component shifts trust rather than eliminating it:
1. **Physical Device Integrity:** You trust your physical devices are not compromised.
2. **Cloudflare Integrity:** You trust Cloudflare not to correlate your WARP tunnel with exit traffic.
3. **Tailscale Integrity:** You trust Tailscale's coordination server not to MITM WireGuard keys.

The goal: no single provider has a complete picture. Your ISP sees an encrypted tunnel. WARP sees traffic with no account identity. Tailscale sees node topology but not content.

9.1 Exit Node IP Rotation (Design Decision)

A common question for privacy architectures is whether the system should aggressively rotate its public exit IP (e.g., every 5 minutes, like a Tor circuit or a proxy rotator).
```

`nullexit` uses Cloudflare WARP via Gluetun. WARP provides **anonymity in a crowd** rather than aggressive rotation. Thousands of users share the same Cloudflare Edge IP simultaneously, making your traffic indistinguishable from the herd. The IP may occasionally rotate on its own (historically surfaced as a `ROTATE` event in `output.log`), but it remains generally stable.

Why aggressive IP rotation was intentionally excluded:
1. **Broken TCP Connections:** Every rotation instantly drops all long-lived TCP sessions (SSH, large downloads, WebSockets, gaming, Zoom).
2. **CAPTCHA & 2FA Hell:** Modern web security (e.g., Cloudflare, Akamai) aggressively flags IP-hopping as bot behavior. Aggressive rotation guarantees the user will be bombarded with "Verify you are human" CAPTCHAs and "New Login Detected" 2FA emails constantly.
3. **WireGuard Architecture:** WireGuard is stateless and does not natively support IP hopping on the fly. To rotate an IP, the VPN client must actively drop the tunnel, request a new endpoint, and re-handshake. This introduces latency gaps where the kill-switch would repeatedly trigger and blackhole the user's internet.

Verdict: For a daily-driver secure gateway, shared-IP crowding (WARP) is superior to aggressive rotation. As an added benefit, Cloudflare WARP IPs are actually highly static (tied to the persistent WireGuard public key generated for your device identity). Because your IP stays static and is part of Cloudflare's own trusted network, it builds a high "trust score," virtually eliminating CAPTCHA loops while still hiding you within the massive crowd of other users in that data center. All mesh devices routing through your gateway will reliably share this exact same trusted IP.

(Note: The only exception where aggressive IP rotation is genuinely required is for specialized offensive security taskslike high-volume web scraping or dark web researchwhere avoiding IP bans or active tracking overrides the need for TCP stability.)

9.2 Cloudflare WARP Privacy & Logging

Because `nullexit` uses Cloudflare WARP (via Gluetun) as its upstream exit node, the system inherits Cloudflare's consumer privacy policy.

What Cloudflare DOES NOT log (Strict No-Logs Policy):
* **Browsing History:** They do not log the domains or URLs you visit.
* **Traffic Destinations:** They do not log the IP addresses of the servers you connect to.
* **Content:** They do not inspect or log the payload/data of your traffic.
* **Data Brokering:** They explicitly guarantee they will never sell or rent your data to advertisers.

What Cloudflare DOES log (Minimal Telemetry):
* **Data Transfer Volume:** Total bandwidth used (required to enforce 10GB/mo quotas if you are not using a paid WARP+ key).
* **Performance Metrics:** Anonymized connection speeds and latency to optimize their network routing.
* **Aggregate Volume:** Total traffic volumes to popular destinations in aggregate (e.g., "datacenter X sent 50 TB to Netflix"), stripped of all user IDs.
* **Device ID (Public Key):** They store the persistent WireGuard public key generated by your container to authorize the connection.

Privacy Verdict:
Cloudflare WARP provides excellent **historical privacy**. Because they do not log your destinations, they cannot comply with historical subpoenas asking what websites you visited yesterday. However, it is **not absolute anonymity**. They still know your real ISP IP address while you are actively connected. For daily-driver privacy against ISPs, hackers, and corporate trackers, it is top-tier. For nation-state evasion, use Tor (see 11).

9.3 OPSEC: Python Runtime Airgap (Isolated Mode)
`nullexit` explicitly relies on the host's native system Python 3 runtime for critical components like `scripts/dns-proxy.py`. Because this script runs with elevated privileges (`sudo -n`) to bind to the privileged UDP/53 socket, modifying or polluting this Python environment is strictly forbidden.

The Zero-Dependency Rule & Isolated Mode (`-I`)
You must **NEVER** install third-party `pip` packages (e.g., `requests`, `pyyaml`) into the system Python environment that `nullexit` uses.
- All Python scripts in this repository must be written using *only* the Python Standard Library (`socket`, `urllib`, `json`, `ssl`).
- Third-party packages introduce a massive supply-chain attack vector. A compromised `pip` package installed globally could allow local malware to hook into the Python runtime and exploit the `sudo` privilege of the DNS proxy to gain full root access.
- To enforce strict immunity against user-installed malicious packages, `toggle.sh` explicitly invokes the proxy using **Python's Isolated Mode** (`python3 -I`). This flag forces the interpreter to completely ignore `PYTHONPATH`, `PYTHONHOME`, and the user's `site-packages` directory, executing exclusively from the read-only, Apple-signed system framework.

9.4 Recommended Upgrades

Layer	Default	Upgraded
VPN Exit	Cloudflare WARP (US)	Mullvad (Swedish, proven no-logs, anonymous payment)
DNS Resolver	AdGuard via WARP	AdGuard via Mullvad DoH (Cloudflare-blind)
Coordination	Tailscale (US, OAuth)	Headscale (self-hosted, no third-party)
Auth	OAuth / persistent keys	Ephemeral auth keys (no identity persistence)
Content	WireGuard asymmetric	WireGuard + PSKs (coordination server blind)

```
#### Mullvad (Replacing WARP)
Swedish-based. Police raided their offices in 2023 and left empty-handed  no logs exist to seize. Accounts are random numbers. Payment by cash or Monero. Replace `VPN_SERVICE_PROVIDER=custom` with `VPN_SERVICE_PROVIDER=mullvad` in `docker-compose.yml`.

#### Mullvad DoH Upstreams
Point AdGuard's upstream resolvers to `https://adblock.dns.mullvad.net/dns-query`. Cloudflare WARP only sees encrypted HTTPS to Mullvad's IPs  completely blind to your DNS queries.

#### Headscale
Self-hosted Tailscale coordination server. Eliminates Tailscale the company, keeping all node topology under your control.

#### Ephemeral Auth Keys
Generate in the Tailscale admin panel. Used once to authenticate; node disappears from the admin panel after disconnect. Breaks persistent identity linkage.

#### Pre-Shared Keys (PSKs)
Add a WireGuard PSK between peer nodes. Adds a symmetric encryption layer that Tailscale's coordination server has no knowledge of, blinding it from reading transit content.

### 9.5 Structural Limitations
- **Traffic Analysis / Metadata:** Passive fiber tapping (UPSTREAM) can observe timestamps, data volumes, and IP ranges without reading content. Defeating traffic analysis requires mixing networks (Tor) or continuous dummy packet padding  both heavily degrade performance.
- **Targeted State-Level Adversaries:** Consumer-grade VPNs are not a complete shield against a nation-state actively targeting you. This stack is designed to defeat bulk passive surveillance and corporate dragnet collection.

---

## 10. Per-Device Access Control

Because every mesh device routes internet-bound traffic through the `warp` container's `FORWARD` chain, they all appear with their stable `100.x.x.x` Tailscale IP as the `src` address. This gives two powerful enforcement surfaces:

### IP & Port Level (`scripts/routing-fix.sh`)
Write raw `iptables` rules targeting specific devices. The 30-second idempotent loop auto-survives Gluetun reconnects:

```bash
# Block a specific device from a target subnet
iptables -I FORWARD -s 100.x.x.x -d 203.0.113.0/24 -j DROP

# Restrict a device to only HTTP/HTTPS
iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 3000 -j DROP
iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 443 -j DROP

# Block a device from internet egress entirely (mesh only)
iptables -I FORWARD -s 100.x.x.x -j DROP

# Time-based (e.g., cut off 10 PM  7 AM)
iptables -I FORWARD -s 100.x.x.x -m time --timestart 22:00 --timestop 07:00 -j DROP
```

DNS Level (AdGuard REST API)
AdGuard Home supports per-client rules keyed by Tailscale IP:

```bash
curl -X POST http://localhost:80/control/clients/add \
  -H "Content-Type: application/json" \
  -d '{
    "name": "kids-ipad",
    "ids": ["100.x.x.x"],
    "blocked_services": ["youtube", "tiktok", "instagram"],
    "safesearch_enabled": true
  }'
```

---

11. Tor & OPSEC Architecture

11.1 Tor-over-WARP Topology

When extreme anonymity is required, integrating the Tor network into `nullexit` provides a powerful Tor-over-VPN topology:
* **The University/ISP** sees only an encrypted Cloudflare WireGuard tunnel, rendering them completely blind to the fact that you are using Tor (effectively bypassing Enterprise firewall Tor blocks).
```

```
* **Cloudflare** sees an encrypted TCP stream heading to a Tor Guard Node. They cannot see your DNS lookups, the destination website, or the content of your traffic.
* **The Tor Exit Node** sees your traffic, but not your real IP (only the Cloudflare IP).

#### The "Transparent Proxy" Trap (What NOT to do)
It is tempting to build a system-wide "Transparent Tor Proxy" that forces all host traffic (e.g., standard Chrome/Safari browsers, background iCloud syncs) through Tor automatically. **This is a massive OPSEC trap and destroys anonymity.**
Modern operating systems and browsers will instantly leak your identity through background syncing (e.g., Apple Mail, Google Drive) and browser fingerprinting (canvas fingerprints, WebRTC, screen resolution). The Tor network protects your *IP*, but if your payload contains identifying cookies or unique fingerprints, the Tor Exit Node (which could be run by malicious actors) will instantly tie your session to your real identity.

#### The nullexit Implementation: Isolated Procedural Proxy
To safely integrate Tor without triggering the transparent proxy trap, `nullexit` implements an **Isolated Tor SOCKS5 Proxy** with **Procedurally Generated Ports**:
1. **Isolated Container:** A lightweight `tor` Docker container runs securely inside the `warp` network namespace. It does not hijack system traffic.
2. **Opt-In Usage:** You must consciously configure a highly-hardened browser (like the official Tor Browser or a dedicated Firefox profile) to use the proxy. This prevents accidental background sync leaks.
3. **Procedurally Generated Ports:** Rather than hardcoding the standard proxy ports (like `127.0.0.1:9050` or `1080`), `toggle.sh` randomizes all internal proxy ports on every boot into the ephemeral range (10000-65000) and silently writes them to `/tmp/nullexit-ports.env`. This completely defeats static fingerprinting by malware.
4. **Kernel-Level Collision Avoidance:** To prevent the randomizer from picking a port already in use by a root daemon or Apple service (which would crash the gateway), the script directly queries the macOS kernel socket table (`netstat -an -p tcp`) to verify the port is absolutely free before assigning it.
5. **The Honey-Port Tripwire:** While procedural ports defeat static fingerprinting, advanced malware might attempt a full sequential port scan of `localhost` to find the hidden proxy. To counter this, `toggle.sh` spawns a "Honey-Port" (a fake random port with a netcat listener attached). If any local application attempts to connect to the Honey-Port, it instantly logs a critical security alert to `output.log` and fires a macOS desktop notification, serving as an early-warning Intrusion Detection System (IDS) for local malware.

> **Threat Model note:** The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). Port randomization and the Honey-Port only defeat blind, generic port-scanners  they do NOT protect against a targeted local process that reads `/tmp/nullexit-ports.env`, greps process memory, or runs `lsof -i`. The proper mitigation is a host-level loopback-filtering firewall (LuLu/Little Snitch). See 15.9 (Threat Model: Local Proxy Discovery) for the full analysis and the rogue-binary verification test.

### 11.2 Hardening: Tor Bridges and obfs4 Pluggable Transports
For users operating under severe Deep Packet Inspection (DPI), `nullexit` supports configuring Tor to connect exclusively through private bridges using the `obfs4` pluggable transport.

When `TOR_USE_BRIDGES=true` is set:
- **What it does:** It hides the entry IP from being matched against the public Tor consensus list, and disguises the traffic's protocol fingerprint from the WARP exit provider (Cloudflare).
- **What it does NOT do:** It does *not* add any extra layers of anonymity beyond what Tor already provides. It is purely a censorship-circumvention and obfuscation mechanism.
- **Limitations:** `obfs4` is highly effective against passive DPI, but it is not guaranteed to defeat active-probing-resistant detection by a highly sophisticated state-level adversary.

> **Architecture Rule:** The `nullexit` gateway deliberately does NOT bundle, hardcode, or automatically fetch bridge lines. Bridge lines are highly sensitive and expire. Users must obtain their own bridge lines from `https://bridges.torproject.org` and manually populate the `tor-bridges.txt` file. If bridges are enabled but the file is missing or empty, the Tor container will intentionally crash and fail to start rather than silently falling back to public guard relays.

**The Bootstrapping Paradox (Out-of-Band Key Exchange)**
Automating the bridge acquisition process (e.g., scripting the gateway to log into the Telegram `@GetBridgesBot` via MTProto API) introduces catastrophic OPSEC flaws:
1. **Identity Leaks:** Baking a Telegram session into the container explicitly links the "anonymous" gateway to a real-world physical phone number.
2. **The Chicken & Egg Deadlock:** In heavily censored regimes, Telegram itself is usually blocked. If the container requires Telegram to fetch bridges to bypass the firewall, it will fail because it cannot bypass the firewall to reach Telegram.
3. **Bridge Exhaustion & CAPTCHAs:** Automated scraping behaves identically to state-sponsored scrapers, triggering Tor's anti-bot CAPTCHAs which leads to silent proxy failures and exhausts the limited public bridge pool.

Instead, users should rely on a less secure layer to manually reach Telegram, obtain the bridges, and populate `tor-bridges.txt`. This can be done via the standard `nullexit` WARP tunnel, a mobile phone on cellular data, Mullvad VPN, or a secure remote VPS. *(Note: We plan on fully integrating Mullvad and VPS architectures directly into `nullexit` in the future, but have not had the time to build it yet).* This intentionally "air-gaps" the cryptographic key exchange from the proxy infrastructure, leaving zero API tokens and zero network traces for an adversary to exploit.

### 11.3 Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification

#### Context
```

```
Unlike the host's SOCKS5 fallback path (whose kill-switch interaction is documented in 15.7), the `tor` container's connection status and fail-closed behavior under VPN tunnel drops is not governed by macOS `pf` rules. Instead, it relies on network namespace inheritance inside Docker Compose.

#### Mechanics
1. **Shared Network Namespace:** The `tor` service is configured with `network_mode: service:warp` in `docker-compose.yml`. This shares the network namespace of the `warp` (Gluetun) container.
2. **Inherited Egress Rules:** All socket communication within the namespace is bound to the same networking stack. Gluetun manages this stack by redirecting traffic through `tun0` (the WireGuard/WARP tunnel) and applying `iptables` rules that explicitly drop outbound traffic originating on `eth0` (the physical/bridge interface), except for permitted local subnets or the VPN server's endpoint.
3. **Tunnel Fail-Closed:** If the WARP connection drops or the `tun0` interface is brought down, the routing table entries for `tun0` become invalid or disappear. Because the `iptables` rules in the shared namespace continue to drop any outgoing internet traffic attempting to exit via `eth0`, the `tor` container cannot leak raw traffic to the host's underlying networks. Egress fails closed.

#### Verification / Manual Test Case
To manually verify that the Tor container fails closed and does not leak traffic when the tunnel dies:

1. **Verify active path:** Ensure the gateway is active (`toggle.sh start`) and fetch the Tor port (`$TOR_SOCKS_PORT` in `/tmp/nullexit-ports.env`). Verify Tor traffic routes correctly:
   ```bash
   curl --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip
   ```
   (This should output a Tor exit node IP).

2. **Simulate tunnel drop:** Bring down the virtual interface inside the shared network namespace:
   ```bash
   docker compose exec warp ip link set dev tun0 down
   ```

3. **Re-run the egress check:**
   ```bash
   curl --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip
   ```
   * **Expected Result:** The request must hang and eventually time out, or immediately fail with a proxy error (e.g., `curl: (97) SOCKS5: connection failed` or `curl: (7) Failed to connect`).
   * **Leaked Egress Check:** Check the host's Wi-Fi router or firewall logs. No outbound packets from the Tor container should route directly over the bridge network to the internet.

11.4 Transparent .onion Routing via RFC 2544 Subnet (198.18.0.0/15)

Context
To enable seamless dark web browsing across the entire mesh, the gateway requires a mechanism to dynamically map `.onion` domains to IP addresses that the host operating system and network layers can route.

IP Address Conflict & Solution
1. **Initial Conflict:** Initially, the Tor virtual address network was configured to map `.onion` sites within the `10.192.0.0/10` range. However, the Colima Docker daemon dynamically allocated `10.200.1.0/24` to the containers. Because this Docker subnet mathematically falls within the `10.192.0.0/10` block, a routing loop occurred: all host packets destined for the Docker containers on ports like `9050` or `9051` were intercepted by the transparent proxy NAT rules and dropped, causing proxy failures.
2. **Tailscale Private IP Exclusion:** Shifting the subnet to another private range like `10.123.0.0/16` resolved the port conflict, but led to connection timeouts. On macOS, Tailscale's Exit Node routing actively excludes RFC 1918 private subnets (e.g. `10.0.0.0/8`, `172.16.0.0/12`, `192.168.0.0/16`) to maintain accessibility to local LAN resources (like printers/routers). Thus, packets targeting `10.123.x.x` were dropped locally by Tailscale before reaching the gateway tunnel.
3. **The Benchmarking Subnet Solution:** To bypass these private IP exclusions without manual static routing tables, Tor's virtual network was changed to **`198.18.0.0/15`** (reserved by RFC 2544 for benchmarking). Since this is not a private RFC 1918 range, Tailscale exit-node routing automatically encapsulates and routes all traffic destined for `198.18.0.0/15` through the tunnel to the gateway.
4. **Transparent NAT Redirection:** The `routing-fix` sidecar applies NAT rules in the shared network namespace:
   ```bash
   iptables -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040
   ```
   This intercepts all incoming TCP traffic targeting the mapped `.onion` range and transparently proxies it through Tor's transparent port (`9040`).

Exit Node Exclusions
The gateway supports the ability to exclude specific countries from being used as exit nodes (e.g. to avoid certain jurisdictions or nodes with poor network latency). This is configured globally via the `TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`, which gets dynamically appended to Tor's `ExcludeExitNodes` and strictly enforced with `StrictNodes 1`.

---

12. Censorship-Resistant Transport & Egress Compartmentalization

12.1 Why this is needed
```

```
WARP can be blocked from deployment countries with strict internet controls. The most likely cause isn't that "encrypted traffic is blocked" in general  it's that WireGuard has a recognizable handshake signature, and `WIREGUARD_ENDPOINT_IP=162.159.192.1` on `WIREGUARD_ENDPOINT_PORT=2408` is a fixed, easily-enumerable target (Cloudflare's known WARP range). That combination is exactly what IP/ASN-based and DPI-based blocking is good at catching.

### 12.2 Important correction: Gluetun's `SHADOWSOCKS=on` is the wrong direction
Gluetun's built-in Shadowsocks option runs a Shadowsocks **server** inside the already-connected VPN tunnel, so other LAN devices can reach out through Gluetun's *existing* connection via the SS protocol. It is not a way to make Gluetun itself connect **outbound** through a Shadowsocks server  Gluetun only speaks OpenVPN or WireGuard as its actual transport. There is no `VPN_TYPE=shadowsocks`. Confirmed against Gluetun's own maintainer/community discussion: people have asked for a "Shadowsocks-client" mode specifically to chain through Gluetun, and the standing answer is "run a separate Shadowsocks client container."

### 12.3 Diagnose before building anything
The right fix depends entirely on what's actually being blocked:
- Point Gluetun at a **different provider/IP/ASN** (any other Gluetun-supported WireGuard/OpenVPN provider, or a self-hosted WireGuard server) and test. If that gets through, the block is Cloudflare/WARP-specific, not a general WireGuard block  **no new component needed**, just swap `VPN_SERVICE_PROVIDER` / the endpoint.
- If WireGuard fails regardless of endpoint, the block is protocol-level (DPI fingerprinting the handshake) proceed to obfuscation.

### 12.4 Path A: Swap WARP for a different Gluetun-native transport
Simplest fix, only applies if 12.3 shows the block is Cloudflare-specific. Change `VPN_SERVICE_PROVIDER` / `VPN_TYPE` / endpoint in `docker-compose.yml` and in the now-centralized `WARP_ENDPOINT_*` values in `.env`. Nothing else in the stack needs to change.

### 12.5 Path B: Add an obfuscation hop
Needed if WireGuard itself is being fingerprinted/blocked. Two ways to slot it in:

**B1  Wrap (recommended):** Run a Shadowsocks (or obfs4/v2ray) client as a new sidecar container. Point Gluetun's `WIREGUARD_ENDPOINT_IP` at `127.0.0.1:<local-port>` instead of Cloudflare directly; the sidecar relays that traffic to a Shadowsocks server you run yourself abroad. Keeps Gluetun's kill-switch, healthchecks, and the entire rest of the stack (Tailscale, AdGuard, ipset, `routing-fix.sh`) untouched, since none of them know the transport changed underneath `warp`.

**B2  Replace:** Swap the `warp` service entirely for a Shadowsocks-client + `tun2socks` container that owns the network namespace instead of Gluetun. More invasive  loses Gluetun's built-in kill-switch/healthcheck, which would need to be rebuilt by hand (see Section 12.9 for the pluggable architecture to support this natively).

**Recommendation: B1.** Smaller surface area, preserves the self-healing architecture already built and documented.

### 12.6 Fingerprinting caveat
Plain Shadowsocks can still be caught by active-probing DPI in more aggressive censorship environments. If plain SS gets blocked too, the next step up is an obfuscation plugin (`v2ray-plugin`, `Cloak`) or a newer protocol designed against active probing (VLESS+Reality, Hysteria2). Test plain SS first  don't over-build before confirming it's needed.

### 12.7 AmneziaWG (The High-Speed Alternative)
If you want to maintain the bare-metal speeds of WireGuard while bypassing DPI (e.g. in Egypt or Russia), the best alternative to Shadowsocks/V2Ray is AmneziaWG. AmneziaWG is a fork of WireGuard that pads the handshake packets with randomized "junk" data, entirely destroying the 148-byte DPI signature that firewalls look for.
- **Pros:** Because it runs entirely as UDP without TCP-over-TCP overhead, it completely avoids "TCP meltdown" latency spikes associated with Shadowsocks/V2Ray wrappers.
- **Cons:** You cannot use Cloudflare WARP. Furthermore, it requires completely ripping Gluetun out of the `nullexit` stack and building a custom Docker container, meaning you lose Gluetun's built-in kill-switches and health-check logic.

### 12.8 Infrastructure Costs & Privacy
To use either Shadowsocks (Path B) or AmneziaWG, you cannot use the free Cloudflare WARP infrastructure, because Cloudflare servers only speak standard WireGuard. The software for both custom protocols is 100% free and open-source, but you must host the remote server infrastructure yourself.
- **Free Tier (Oracle Cloud):** You can host your remote server on Oracle Cloud's "Always Free" tier for $0. However, this routes your highly sensitive, censorship-evading traffic directly through Oraclea massive corporate data broker. If privacy is the core objective, handing all your metadata to Oracle defeats the purpose.
- **Paid Tier (Hetzner / Mullvad):** The recommended path is to rent a standard ~$4/month Virtual Private Server (VPS) in a free country (e.g., Hetzner in Germany, subject to strict EU GDPR privacy laws) and self-host the server. Alternatively, Mullvad VPN (5/month) provides native v2ray/Shadowsocks bridges built into their network, allowing you to bypass censorship with a strict zero-logs policy. It is highly recommended to pay with untrackable payment methods to maintain complete anonymity, which these providers typically support without requiring local residency.

### 12.9 The Future: Egress Compartmentalization (Pluggable Architecture)
To natively support invasive replacement paths like **Path B2** or **AmneziaWG** (Section 12.7) without breaking the core `nullexit` routing and monitoring logic, the architecture must be refactored into a modular, "pluggable" design.
```

```
1. **The Egress Plugin System:** Currently, WARP-specific logic (like `cdn-cgi/trace` health checks and hardcoded interface names) is scattered across `toggle.sh` and `scripts/routing-fix.sh`. This should be abstracted into a `scripts/plugins/` directory, where each egress method (`warp.sh`, `v2ray.sh`) implements standardized functions: `get_egress_interface()`, `check_health()`, and `get_bypass_ips()`.
2. **Dynamic Docker Compose:** Based on an `EGRESS_TYPE` variable in `.env`, dynamic compose file generation would spin up only the required egress containers (e.g., skipping the `warp` container when `EGRESS_TYPE=v2ray` is set).
3. **Agnostic Routing & Watchers:** `routing-fix.sh` would dynamically read the `EGRESS_INTERFACE` from the active plugin to apply `iptables` rules, and the background watchers in `toggle.sh` would become an agnostic `Egress Watcher` that invokes `check_health()` periodically.

This compartmentalization transforms `nullexit` into a universal, censorship-resistant gateway where the entire egress layer can be swapped by changing a single `.env` variable and running `./toggle.sh --restart`.

---

## 13. Roaming & Recovery Design: How nullexit Mimics Commercial VPNs

The `nullexit` roaming/network recovery subsystem replicates the exact engineering mechanics used by commercial, premium VPN clients (like Mullvad or NordVPN) to transition across Wi-Fi networks safely and seamlessly. This is the design overview; the specific bugs encountered and fixed along the way live in the Incident Log (15.5 Roaming/Sleep-Wake, 15.7 Tailscale P2P/SNAT).

### The 4-Step Transition Cycle
When your Mac disconnects from one Wi-Fi and associates with another, `nullexit` coordinates the following events:

1. **Firewall Lock (Kill-Switch)**: The macOS Packet Filter (PF) anchor `com.apple/nullexit` remains locked on the physical interface (`en0`), blocking all direct outbound IPv4 and IPv6 traffic. This guarantees that your real ISP IP or unencrypted DNS requests **never leak** onto the new network while the connection is rebuilding.
2. **System Event Interception**: A launchd LaunchAgent (`watcher.sh`) listens to the macOS System Configuration framework for the `State:/Network/Global/IPv4` key changes, spawning `recover.sh --post-wake` in under 500ms when a change is detected.
3. **Bypass Routing**: The script dynamically resolves the new physical network's IP gateway (e.g. `192.168.137.1`), cleans up stale bypass routes from the previous network, and adds static bypass routes pointing to Cloudflare's WARP endpoints and the Tailscale control plane directly via the new gateway. This allows the VPN client to negotiate its handshake outside of the tunnel.
4. **Hole-Punching / NAT Re-negotiation**: The script runs a target check on the `warp` container's tunnel. If the WireGuard connection is wedged by the network switch, it force-recreates the container to establish a fresh NAT binding to Cloudflare, restoring connection in ~15 seconds.

---

## 14. Deployment, Packaging & Known Limitations

### 14.1 Packaging nullexit as a macOS Application (.dmg)

nullexit can be packaged into a native `.app` bundle, but this introduces nested virtualization requirements.

nullexit uses **Colima** (Apple Virtualization Framework `vz`) to run a Linux VM hosting Docker. Installing the `.app` inside a virtualized macOS environment (Parallels, cloud Mac) means running a Linux VM inside a macOS VM  requiring hardware-level **nested virtualization**.

#### Hardware Support
- **Supported:** Apple Silicon M3/M4 (macOS Sequoia 15+), Intel Macs with VMX passthrough.
- **Unsupported:** M1, M2, A18 Pro. Colima crashes silently; the terminal (`setup.sh`) surfaces crash logs.

#### Build Steps
```bash
# 1. Verify nested virtualization support
sysctl -a | grep hv_nested_virt_supported  # expect: 1

# 2. Compile the AppleScript launcher
osacompile -o "Nullexit.app" "Toggle Gateway.applescript"

# 3. Embed scripts into app resources
mkdir -p "Nullexit.app/Contents/Resources/scripts"
cp toggle.sh setup.sh recover.sh docker-compose.yml "Nullexit.app/Contents/Resources/"

# 4. Package into DMG
hdiutil create -volname "Nullexit Installer" -srcfolder "./Nullexit.app" -ov -format UDZO Nullexit.dmg

# 5. Bypass Gatekeeper (unsigned app)
xattr -cr /Applications/Nullexit.app
```

Without a paid Apple Developer certificate ($99/yr), users see an "App is damaged" Gatekeeper warning and need to run step 5.

14.2 Moonlight/Sunshine + Tailscale Race Condition
```

```
630
631 When streaming from a remote Windows host running Sunshine over a Tailscale mesh (e.g., while the client is on
        a cellular hotspot), a race condition can occur on boot:
632
633 1. Both Tailscale and Sunshine are set to start automatically on boot.
634 2. Tailscale takes a few seconds to fully initialize its `100.x.x.x` virtual adapter and establish a connection
        .
635 3. **Sunshine starts too fast.** It launches before Tailscale is fully ready. Since Sunshine only binds to
        network interfaces that are active at the exact moment of startup, it misses the Tailscale adapter entirely
        .
636 4. Moonlight clients (even when configured with the correct Tailscale IP) will fail to connect because Sunshine
         isn't listening on that interface.
637
638 **The "Magic Toggle" Symptom:**
639 If the user manually enables or disables Wi-Fi on the host PC, it triggers a global Windows network change
        event. Sunshine listens for these events, re-enumerates all network adapters, and says, "Oh, there's a
        Tailscale adapter here!" and finally binds to it. Suddenly, Moonlight can connect.
640
641 **The Permanent Fix:**
642 On the Windows host, open `services.msc`, locate the **Sunshine Service**, and change its Startup type from **
        Automatic** to **Automatic (Delayed Start)**. This forces Windows to wait a minute or two after booting
        before launching Sunshine, ensuring Tailscale's virtual adapter is fully initialized first.
643
644 ### 14.3 Chrome Remote Desktop Performance (UDP NAT)
645
646 #### Observation (July 11, 2026)
647 When the user connects to their host Mac using Chrome Remote Desktop (CRD) while `nullexit` is active, the
        connection suffers from massive latency, lag, and degraded video quality. AdGuard Home's `blocked.log`
        shows zero blocked DNS requests for Google, WebRTC, or `chromoting` endpoints, indicating this is not a DNS
         sinkhole issue. *(See also 15.11 for the separate CRD-on-remote-device connection-failure quirk.)*
648
649 #### Root Cause
650 This is a fundamental limitation of tunneling all host traffic through a strict commercial NAT (Cloudflare WARP
        ).
651 1. Chrome Remote Desktop attempts to use **STUN (UDP hole-punching)** to establish a lightning-fast, direct
        peer-to-peer WebRTC connection between the client device and the host Mac.
652 2. Because all non-Tailscale traffic is forcefully routed through the `warp` container, the STUN packets exit
        the network from a Cloudflare IP.
653 3. Cloudflare's enterprise NAT infrastructure strictly drops unsolicited inbound UDP packets. The UDP hole-
        punch fails.
654 4. Because a direct P2P connection cannot be established, CRD falls back to **Relay Mode** (TURN), bouncing all
         video data back and forth through Google's cloud servers instead of sending it directly over the local or
        mesh network. This relaying introduces massive latency.
655
656 #### Workaround / Fix
657 This lag is the accepted "tax" for maintaining absolute cryptographic anonymity; the only way to restore CRD
        speed would be to leak its UDP traffic outside the tunnel (violating zero-trust).
658 To achieve low-latency remote access, users should bypass CRD entirely and use macOS's built-in **Screen
        Sharing (VNC)** directly over the Tailscale IP. Tailscale's sophisticated custom DERP/WireGuard
        architecture successfully UDP hole-punches through strict NATs, providing a direct, encrypted, high-speed
        connection without relying on external cloud relays.
659
660 ### 14.4 Dynamic IP Fetch vs. MagicDNS
661
662 #### Observation
663 The `nullexit` gateway uses a dynamically fetched IP address (cached in `.gateway_ip`) to configure host
        routing and the `pf` kill-switch, rather than utilizing Tailscale's user-friendly MagicDNS hostname (e.g.,
        `nullexit-gateway`).
664
665 #### Why MagicDNS is Not Used
666 1. **The "Chicken and Egg" Bootstrapping Problem:**
667    macOS's Packet Filter (`pf`) and low-level routing commands (`route add`) operate strictly at Layer 3 and
        require raw IP addresses. If `toggle.sh` or `recover.sh` attempted to use `nullexit-gateway`, the host Mac
        would need to perform a DNS lookup to resolve it. However, the very first step of the gateway's
        initialization is to completely hijack the host's DNS and point it *at the gateway itself*. The Mac cannot
        resolve the gateway's MagicDNS name if it needs the gateway's IP to know where to send the DNS query!
668 2. **Dynamic Self-Healing (Avoiding Stale Cache):**
669    To solve the bootstrapping problem without causing bugs if the node is deleted and re-authenticated on the
        Tailscale Admin Console, `toggle.sh` explicitly runs `docker compose exec ... tailscale ip -4` during boot
        to fetch the absolute freshest IP dynamically. It saves this to `.gateway_ip` so that fast-roaming scripts
        like `recover.sh` can instantly retrieve the routing target without incurring the latency of querying
        Docker.
670
671 ### 14.5 Battery & Power Measurement Quirks
672
673 When trying to optimize this framework (or any daemon) for battery life, developers often look for a way to
        programmatically measure the exact Wattage consumed by a specific process (e.g., "How many Watts is `toggle
        .sh` using?").
674
675 **This is a hardware impossibility on macOS and Linux.**
676
```

Why Activity Monitor is Lying to You
Your machine's logic board only has physical power sensors for the *entire* CPU package, the GPU, and the RAM. It knows the CPU is pulling 5W total, but it has no physical hardware capability to know whether Docker is using 3W and Chrome is using 2W.

The "Energy Impact" score you see in macOS Activity Monitor is a **synthetic, heuristic score** calculated by `powerd`. It penalizes apps for waking up the CPU (idle wakes), doing disk I/O, or keeping the screen awake. It is a relative score, not actual Wattage.

The Proper A/B Testing Methodology
If you want to truly know how many Watt-hours or mAh your background daemons (`routing-fix.sh`, the DNS/WARP watchers, etc.) are costing you, you must perform a hardware-level A/B drain test:

1. Unplug the laptop, run the gateway, and note your exact battery mAh using:
   ```bash
   ioreg -l | grep "AppleRawCurrentCapacity"
   ```
2. Leave the laptop idle for exactly 1 hour.
3. Run the command again to see exactly how many mAh were drained.
4. Turn the gateway off and repeat the test for another hour.

The delta in mAh drained between the two tests is the true, exact cost of running the software. When optimizing polling loops (e.g., changing `sleep 5` to `sleep 30`), this is the only reliable way to measure the actual battery life returned to the user.

14.6 Host Lockdown Mode (Why It's Impossible on macOS)

A common privacy feature request is a "Host Lockdown Mode" (Mode 3): A mode where the Docker exit node remains perfectly functional for remote devices, but the host machine itself is completely blocked from accessing the internet (to prevent even encrypted host traffic from leaking metadata).

Why this is impossible on macOS:
On Linux (e.g., a Raspberry Pi), Docker runs natively. The host's traffic traverses the `OUTPUT` iptables chain, while Docker traffic traverses the `FORWARD` chain. We could easily drop all `OUTPUT` traffic to kill the host's internet, and Docker would continue routing traffic perfectly.

However, Docker on macOS runs inside a Linux Virtual Machine (via `qemu` or Apple `Virtualization.framework`). This VM runs as a standard macOS user-space application. If you configure the macOS firewall (`pf`) to drop all host outbound traffic, it will indiscriminately drop the VM application's traffic as well, instantly killing the Cloudflare WARP tunnel and the exit node.

Writing complex macOS `pf` rules to "allow the VM app, allow Cloudflare WARP IPs, allow Tailscale DERP IPs, but block everything else" is highly fragile and heavily discouraged. Since the current `toggle.sh` default mode already fully encrypts the host's traffic via the WARP tunnel, the host is already protected from ISP leaks.

Part III Incident Log & Resolved Issues

This is the single, deduplicated log of every incident, bug, and resolved issue. Entries are organized into subsystem groups (15.115.12). Each entry follows a **Symptom / Root Cause / Fix** shape where applicable; packet-level analyses are preserved as clearly-labeled **Deep dive** blocks. The terse dated changelog lives in 17; this section holds the full post-mortems.

15. Incident Log

15.1 Exit-Node Return Path & SOCKS5 Failover

15.1.1 Exit Node Return Path (`rx 0`) & The SOCKS5 Failover RESOLVED
Symptom: Gateway showed `tx 936 rx 0` transmitted but never received return traffic. For days, Tailscale's exit node refused to return traffic back to the client, leading to the assumption that Tailscale's userspace forwarding was bypassing `iptables` and ignoring the WARP `tun0` interface. In a desperate attempt to fix this, a **SOCKS5 proxy** (`scripts/socks5-proxy.py`) was introduced as a replacement.

Root Cause: Tailscale's userspace forwarding was *not* the problem. `docker-compose.yml` included `FIREWALL_OUTBOUND_SUBNETS=100.64.0.0/10`. Gluetun created a strict routing rule (`ip rule add to 100.64.0.0/10 lookup 199`, priority 99) that forced all Tailscale CGNAT traffic to bypass the VPN via the unencrypted Docker bridge (`eth0`). Return packets were sent to the macOS host instead of back through `tailscale0`, blackholing all return packets back to the client.

Fix: Removed `FIREWALL_OUTBOUND_SUBNETS` entirely, immediately restoring the Tailscale exit node. `scripts/routing-fix.sh` now injects `lookup 52`, and return packets correctly flow back into `tailscale0`. Instead of removing the SOCKS5 proxy, we repurposed it into a **bulletproof failover** for when restrictive Wi-Fi networks block Tailscale UDP ports.

Architecture after the fix:
```
PRIMARY (Exit Node):
DNS/TCP/UDP: Apps  Tailscale Exit Node  gateway container  tun0  WARP  internet
```

```
FAILOVER (SOCKS5 - if Tailscale is blocked by local Wi-Fi):
DNS: Browser  127.0.0.1:53  Python proxy  TCP:5354  AdGuard  WARP
TCP: Apps  SOCKS5:1080  gateway container  tun0  WARP  internet
Mesh: SSH/ping 100.x.x.x  utun5  WireGuard  peers (independent of proxy)
```

Deep dive: Exit Node Return Path Analysis (June 26, 2026).
This preserves the detailed packet-level analysis that led to identifying and fixing the `rx 0` bug.

The exit node was probably never going to work in this container topology. Here's the packet path traced by reading `ip rule show` and the routing tables live:

Outbound (Phone-1 internet):
1. Phone-1 sends packet to gateway via Tailscale mesh
2. Gateway's tailscale0 (kernel TUN, `TS_USERSPACE=false`) decrypts packet enters FORWARD chain
3. Packet has src=`100.87.42.87` (Phone-1), dst=some internet IP
4. FORWARD chain (nftables): Rule 1 matches (`-i tailscale0 -j ACCEPT`)
5. Routing decision for the forwarded packet. Which ip rule matches?
   - Rule 98: `to 172.18.0.0/16` no (dst is internet)
   - Rule 99: `to 100.64.0.0/10` no (dst is internet)
   - Rule 100: `from 172.18.0.2` **NO** (src is `100.87.42.87`, not the container IP!)
   - Rule 101: `not fwmark 0xca6c lookup 51820` **YES** (no fwmark on forwarded packets)
6. Table 51820: `default dev tun0` packet goes out through WARP
7. NAT POSTROUTING: `MASQUERADE on tun0` src becomes WARP IP
8. Packet exits through WARP to internet (in theory)

Return (internet Phone-1):
1. Return packet arrives on tun0, conntrack de-MASQUERADEs dst back to `100.87.42.87`
2. Routing decision for dst=`100.87.42.87`:
   - Rule 99: `to 100.64.0.0/10 lookup 199` **YES**
3. Table 199: `100.64.0.0/10 via 172.18.0.1 dev eth0` **sends it to the Docker gateway!**
4. **This is wrong.** The packet should go to tailscale0 (to be re-encrypted and sent to Phone-1), but table 199 sends it out eth0 to the Docker bridge host.

Table 199 is the smoking gun. It has one route: `100.64.0.0/10 via 172.18.0.1 dev eth0`. This sends ALL Tailscale CGNAT return traffic out the Docker bridge to the host not back through tailscale0. The return packet leaves the container's network namespace entirely, hits the host's routing stack, and gets dropped because the host has no idea what to do with a raw packet for `100.87.42.87`.

Why table 199 exists and why it's wrong for exit node traffic. Table 199 + the CGNAT rule (pref 99) was designed so that the **container itself** can reach other Tailscale peers (e.g., for mesh management, DERP relay). For that use case, routing CGNAT traffic via the Docker bridge to the host (which has its own tailscaled) makes sense. But for **forwarded exit node traffic**, the return packet needs to go back through tailscale0 inside the container not out to the host. The CGNAT rule intercepts the return packet before it can be routed back to tailscale0 through the main table or Tailscale's own routing (table 52, which is at pref 5270).

The SOCKS5 path works because it dodges all of this. SOCKS5 proxy creates connections **from the container's own IP** (`172.18.0.2`). So:
- ip rule 100 (`from 172.18.0.2 lookup 200`) matches
- Table 200 has `default dev tun0`
- Return traffic comes back to `172.18.0.2` on tun0, conntrack de-NATs, proxy sends response to client

No FORWARD chain involved. No CGNAT rule involved. No table 199. It just works.

Summary of findings:

Finding	Status	Evidence
Table 199 CGNAT route hijacks exit node return traffic	**Fixed**	Changed `lookup 199` to `lookup 52` in `scripts/routing-fix.sh`. Return packets now route via tailscale0.
FORWARD RELATED,ESTABLISHED missing	**Fixed**	Installed `iptables` via apk in `docker-compose.yml`. Rule now injects successfully.
DOCKER_GW variable parsing error	**Fixed**	Added `head -1` in `scripts/routing-fix.sh`.
`FIREWALL_OUTBOUND_SUBNETS` was the root cause	**Fixed**	Removed from `docker-compose.yml`. Gluetun no longer hijacks CGNAT routing.
Host tailscaled error state from aggressive `tailscale down`	**Mitigated**	Toggle script now detects and auto-restarts.

15.1.2 Container Default Route Bypasses WARP
Problem: Inside the WARP container's network namespace, the main routing table's default route was via `eth0` (Docker bridge), not `tun0` (WARP tunnel). While gluetun uses policy routing (rule 101 table 51820 `tun0`) for most traffic, traffic originating from the container's own IP (`172.18.0.2`) hits rule 100 (`from 172.18.0.2 lookup 200`), whose default was also via `eth0`. This caused the SOCKS5 proxy's outbound connections to bypass WARP.

Fix: Updated the `routing-fix` container to:
1. Change the main table's default route to `default dev tun0`
2. Change table 200's default route to `default dev tun0`
3. Add a specific route for the WARP WireGuard endpoint (`162.159.192.1`) through `eth0` via the Docker gateway in table 200, preventing a tunnel loop

4. Re-assert all routes every 30 seconds (gluetun may reset them on health checks)
5. Dynamically detect the Docker subnet from `ip route show dev eth0` instead of hardcoding `172.18.0.0/16`

15.1.3 Hardcoded Docker Subnet in Routing Fix
Problem: The `routing-fix` container hardcoded `172.18.0.0/16` and `172.18.0.1` for the Docker bridge subnet and gateway. If Docker's bridge network changed (after `docker network prune`, Colima restart, or on a different machine), these routes would be wrong.

Fix: Detection is now dynamic using `ip route show`:
- `DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')`
- `DOCKER_GW=$(ip route show default | awk '{print $3}')`

This extracts the actual subnet and gateway directly from the routing table, working with any subnet size (`/16`, `/24`, `/20`, etc.).

15.2 Routing Loops & Subnet Collisions

15.2.1 The Infinite Recursive Routing Loop (Hotspot Paradox)
Problem: A critical network topology crash occurs if the host Mac connects to a Mobile Hotspot (e.g., a Windows PC or Phone) that is *also* connected to the Tailscale mesh and has "Use Exit Node" enabled. Because the Mac is hosting the Exit Node, it routes its internet traffic to the Hotspot. The Hotspot intercepts the Mac's traffic on its way to the cellular tower and routes it *back* to the Exit Node (the Mac) via Tailscale. The Mac receives it, tries to send it out again, and the Hotspot intercepts it again. This creates an infinite, recursive encryption loop that instantly destroys network bandwidth and drops all connections.

Fix: This is a physical routing paradox that cannot be resolved in software. The upstream physical router providing internet to the Mac **must** be allowed to route its traffic directly to the physical WAN interface. The user must manually disable "Use Exit Node" on the upstream device providing the hotspot.

Advanced Workaround: If the upstream hotspot is a Windows machine and the user explicitly wants it to remain connected to the Exit Node, a permanent static route can be injected into the Windows routing kernel to forcefully bypass the Tailscale WFP interception for WARP packets. By binding the route to the physical adapter's Interface Index (IF), it becomes a permanent, roaming-aware bypass.

On Windows (Admin Command Prompt):
```cmd
route print  (find Interface List, note the IF number for the active Wi-Fi adapter, e.g. 15)
route -p add 162.159.192.1 mask 255.255.255.255 0.0.0.0 IF 15
route -p add 162.159.193.1 mask 255.255.255.255 0.0.0.0 IF 15
```

On Linux (Root Terminal):
```bash
ip route add 162.159.192.1 dev wlan0
ip route add 162.159.193.1 dev wlan0
```

On macOS (Root Terminal):
```bash
route add -host 162.159.192.1 -interface en0
route add -host 162.159.193.1 -interface en0
```

(Note: If the upstream router is an iPhone or Android device, you cannot inject static routes without jailbreak/root. You must simply disable "Use Exit Node" in the mobile Tailscale app).

This punches a tiny hole straight through the paradox, letting the Mac's WARP packets escape to the physical internet while keeping the rest of the upstream router's traffic securely trapped in the Tailscale Exit Node.

15.2.2 Infinite Egress Routing Loops & Subnet Takeover Paradoxes
Problem: Two distinct network loops can break host egress connectivity entirely when the exit node is active:

1. **The Recursive Host-VM Tunnel Loop:**
   - **Mechanism:** The host Mac's default route points to the Tailscale exit node (`utun*`). The exit node container sends its encrypted tunnel packets (destined for the WARP endpoint `162.159.192.1`) back to the host via the VM NAT. Without a static route exception on the host Mac, the host Mac routes those packets right back into the exit node (`utun*`), creating an infinite loop that blackouts all host network egress.
   - **ARP-Scoping Pitfall:** Simply binding the bypass route to the interface (`route add -host 162.159.192.1 -interface en0`) fails when the default route is deleted or redirected to `utun*`. Since `162.159.192.1` is not local to the physical link, macOS fails to resolve it via ARP, dropping all tunnel packets.
   - **Fix:** Dynamically resolve the active physical gateway IP (e.g. `172.17.0.1`) on startup, and add the static host routes for the WARP endpoints via the gateway IP directly (`route add -host 162.159.192.1 <gateway_ip>`). Additionally, because standalone `tailscaled` on macOS does not automatically modify the system default route via the CLI, we manually redirect the default gateway to `utun*` using `setup_exit_node_routing`.

2. **The Docker Compose Subnet Takeover Collision:**

```
841    - **Mechanism:** To avoid collisions with the host's campus network, Colima's default bridge (`docker.bip`)
         is moved (e.g. to `10.200.0.1/24`). However, because `172.17.0.0/16` is now free inside the VM, Docker
         Compose's IPAM dynamically takes it over for the project's custom network (`nullexit_default`). Because the
          containers receive IPs on `172.17.0.0/16`, the host Mac cannot send local peer discovery packets (e.g.,
         Tailscale `magicsock` packets on `172.17.0.2:41641`) directly. The host routing table sends them out `en0`
         to the physical Wi-Fi, returning `host is down` or `no route to host`.
842    - **Fix:** Explicitly configure a custom default network subnet `10.200.1.0/24` in `docker-compose.yml` to
         prevent Docker Compose from ever reclaiming the conflicting `172.17/172.18` ranges.
843
844 > See also 15.2.1 (The Hotspot Paradox) for the related infinite loop that occurs when the physical upstream
         router is also a Tailscale exit-node client.
845
846 #### 15.2.3 Docker Default Bridge vs Host Wi-Fi Subnet Collision (The 172.17.0.0/16 Subnet Overlap)
847 **Context:** When attempting to ping a local Windows client (`172.17.22.187`) or the default gateway
         (`172.17.0.1`) directly over Wi-Fi, the Mac host returns `No route to host` (displaying a `REJECT` / `!`
         flag in the routing table). Tailscale on the client also drops offline shortly after starting when using
         the exit node, failing to establish direct P2P connections and falling back to slow DERP relays.
848
849 **Root Cause & Subnet Overlap Mechanism:**
850 1. **Docker Default Subnet:** By default, the Docker daemon initializes its default bridge network (`docker0`)
         on the **`172.17.0.0/16`** IP range.
851 2. **Subnet Conflict:** The host's physical Wi-Fi network happens to be assigned the exact same
         **`172.17.0.0/16`** range by the local network router.
852 3. **Routing Clash:** Colima launches a Linux VM on the Mac host to run the Docker engine. Because Docker
         inside that VM creates `docker0` and claims the `172.17.0.0/16` subnet, any packet sent to `172.17.x.x`
         from within the containers (such as Tailscale local discovery) is captured by the VM's internal virtual
         bridge interface (which is down/empty) and is blackholed. On the Mac host, macOS dynamically marks routes
         as `REJECT` (`!`) when local ARP resolution fails, causing local pings to immediately abort with `No route
         to host`.
853 4. **Tailscale Relay Saturation**: Because local peer-to-peer discovery is blackholed, Tailscale falls back to
         public DERP relay servers (e.g. `relay "tor"` in Toronto). When exit-node routing is enabled, all client
         web traffic is routed through the relay. Public relays impose strict rate limits and throttle high-volume
         traffic, resulting in packet drops. This saturation drops Tailscale's control packets, causing the
         connection to time out and showing the client as offline.
854
855 **Fix:**
856 The canonical script for this is `scripts/fix-docker-bridge-collision.sh`. **One command applies the entire fix
         :**
857
858 ```bash
859 bash scripts/fix-docker-bridge-collision.sh        # apply
860 bash scripts/fix-docker-bridge-collision.sh --dry  # preview changes without applying
861 ```
862
863 It auto-detects the host's actual Wi-Fi subnet (no more guessing whether the campus DHCP handed out `172.x`,
         `10.x`, or `192.168.x`), picks a non-conflicting `docker.bip` from a small candidate set (defaulting to
         `172.26.0.1/24` if no overlap is matched), backs up `~/.colima/default/colima.yaml` with an ISO-8601
         timestamp, edits the file **in place** (idempotent  replaces an existing `docker.bip` if present, appends
         under an existing `docker:` block, or creates a fresh block), restarts Colima, rebinds Wi-Fi DHCP, verifies
          the new default route is no longer Docker's bridge, and re-runs both `./toggle.sh` and `scripts/diagnose-
         host-leak.sh` to print the post-fix verdict.
864
865 For reference, the underlying manual fix is: edit the Colima configuration file (`~/.colima/default/colima.yaml
         `) on your Mac to define:
866 ```yaml
867 docker:
868   bip: 172.26.0.1/24   # any /24 that does NOT overlap your Wi-Fi subnet
869 ```
870 Then, restart Colima to apply the new subnet inside the VM:
871 ```bash
872 colima restart
873 ```
874 This forces Docker to use `172.26.x.x` for its default bridge, freeing the physical `172.17.x.x` range and
         letting pings and direct Tailscale peer-to-peer connections flow normally.
875
876 #### 15.2.4 July 4, 2026: The "Network Unreachable" Host Leak & Post-Wake Container Destructions
877 **Symptom:**
878 1. Running `docker compose config` with dummy keys corrupted the user's `.env` file by appending invalid
         WireGuard keys. This caused the `warp` container to become unhealthy on boot, which triggered a full
         teardown of the gateway by `toggle.sh`.
879 2. After fixing `.env`, `toggle.sh` successfully booted the gateway, but the host's internet started bypassing
         the tunnel entirely (a massive host leak). The Cloudflare trace returned `warp=off` and exposed the host's
         physical ISP IP.
880 3. Running `recover.sh --post-wake` would violently force-recreate all gateway containers and drop the host's
         internet connection.
881
882 **Root Cause:**
883 - **Bug 1 (The Host Leak):** In `toggle.sh`, the logic to find the Tailscale `utun` interface used `ifconfig |
         grep -B4 "inet 100."`. Due to the 4 lines of before-context, this regex was accidentally capturing the
         interface *above* the actual Tailscale interface (e.g. grabbing `utun4` instead of `utun5`). Because `utun4
         ` had no IP address, the subsequent `sudo route add -net 0.0.0.0/1 -interface utun4` silently failed with "
```

```
    Network is unreachable", leaving the host's default route exposed.
- **Bug 2 (The Post-Wake Destructions):** The health check in `recover.sh --post-wake` relied on `docker compose exec -T warp curl -sf https://www.cloudflare.com/cdn-cgi/trace`. However, the newer `qmcgaw/gluetun` Alpine-based Docker image no longer ships with `curl` pre-installed. The health check failed with `executable file not found in $PATH`, tricking `recover.sh` into thinking the tunnel was completely dead and needed to be forcefully recreated via `docker compose up -d --force-recreate`.
- **Bug 3 (Restart Flags):** Previous refactors incorrectly permitted `toggle.sh restart` instead of strictly enforcing `toggle.sh --restart`.

**Resolution:**
- Cleaned the `.env` file of duplicate dummy entries.
- Replaced the brittle `grep -B4` interface parsing in `toggle.sh` with a strict AWK parser: `awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':'`. This mathematically isolates the exact interface possessing the Tailscale `100.x.x.x` IP address, preventing the `0.0.0.0/1` route from silently failing.
- Changed the health check in `recover.sh` to use `wget -qO- --timeout=3` which is natively installed in the Gluetun image. This stopped the unnecessary destructive recreations of the containers during post-wake events.
- Reverted the `restart` alias in `toggle.sh` to strictly require `--restart`.

#### 15.2.5 July 4-5, 2026: Tailscale Data-Plane Loop (The DERP Relay Deadlock)
**Symptom:**
After bypassing the Tailscale control plane (`192.200.0.0/16`) to fix the "offline" mesh status, the Mac appeared online. However, external peer devices (like the user's phone on cellular) could not successfully ping or SSH into the Mac. `tailscale netcheck` reported `Nearest DERP: unknown (no response to latency probes)`.

**Root Cause:**
While bypassing the control plane kept the daemon connected to the coordination servers, Tailscale's actual peer-to-peer data plane relies on STUN (for NAT traversal) and DERP relays. These relays span ~80 dynamically allocated IP addresses across multiple cloud providers. Because we were forcing `0.0.0.0/1` through the `utun*` exit node, the Mac's own outbound connection attempts to DERP relays were being routed back into the tunnel. To prevent an infinite loop, Tailscale's daemon dropped its own packets when it saw them returning on the TUN interface. This effectively left the daemon deaf and blind to peer connections.

**Resolution:**
- Modified `add_warp_bypass_routes` in `scripts/common.sh` to execute a live API fetch (`curl -s https://login.tailscale.com/derpmap/default`) during startup.
- The script parses out all active DERP relay IPv4 addresses (~80 IPs) and adds a physical host bypass route for every single one of them.
- These IPs are written to a temporary file (`/tmp/nullexit-derp-ips.txt`) so they can be cleanly un-routed by `remove_warp_bypass_routes` when the gateway shuts down. Peer connectivity (ping, SSH, SFTP) was fully restored.

#### 15.2.6 July 5, 2026: Hardcoded WARP Endpoints in docker-compose
**Symptom:**
The bash scripts allowed overriding the default Cloudflare WARP IP endpoints via `.env` variables (`WARP_ENDPOINT_1` and `WARP_ENDPOINT_2`) to establish bypass routes. However, `docker-compose.yml` statically hardcoded `162.159.192.1` for the `warp` container's `WIREGUARD_ENDPOINT_IP`.

**Root Cause & Risk:**
If a user set a custom endpoint in `.env`, the script would successfully bypass the custom IP on the host level. However, Gluetun would still attempt to connect to the hardcoded `162.159.192.1`. Since `162.159.192.1` was no longer explicitly bypassed, its packets would be sucked into the `utun*` exit node, causing an infinite WireGuard routing loop that instantly breaks the gateway.

**Resolution:**
Modified `docker-compose.yml` to dynamically read the environment variable via `${WARP_ENDPOINT_1:-162.159.192.1}`, ensuring both the host routing scripts and the container use the exact same endpoint.

#### 15.2.7 Incident Post-Mortem: WARP Watcher Race vs Post-Wake (July 10, 2026)
**Symptom:** After switching Wi-Fi networks, the host IP appeared as the raw ISP IP (`132.x.x.x`) instead of a Cloudflare/WARP IP. The gateway appeared to complete post-wake recovery successfully in the logs, but nuclear shutdown fired immediately after and killed everything.

**Root Cause:** A fatal race condition between two concurrent processes:

1. **Wi-Fi roam**  `watcher.sh` fires `recover.sh --post-wake`
2. **Post-wake** finds the `warp` container missing/dead  calls `docker compose up -d --force-recreate` (~35 seconds)
3. **In parallel**, the WARP Watcher (running as a child of `toggle.sh`) polls the warp container every 5s during failures
4. While the container is being recreated it returns `warp=off` for those 35 seconds  7 consecutive failures **WARP SHUTDOWN fires**, calling nuclear `recover.sh`
5. Nuclear `recover.sh` tears down Tailscale, resets DNS to 1.1.1.1, stops all containers  **gateway dead, IP exposed**
6. The post-wake `docker compose up` finishes at ~35s mark, container healthy  but the gateway is already gone

The evidence in `output.log`:
```
[2026-07-10T06:07:27Z] NET: ... bash recover.sh --post-wake
```
```

```
[2026-07-10T06:07:47Z] WARP DOWN  cdn-cgi/trace reports warp=off    watcher starts counting
...
add net 0.0.0.0: gateway utun5    post-wake successfully re-routes at 02:08:22
...
[2026-07-10T06:09:04Z] WARP SHUTDOWN  6 consecutive failures         watcher fires nuclear
```

Fix (July 10, 2026): Added a **WARP Watcher inhibit marker** (`/tmp/nullexit-warp-inhibit.marker`):
- `recover.sh --post-wake` writes the marker at startup **before** any container operations
- The WARP Watcher loop checks for the marker at the top of every iteration; if present it resets `consec_off=0` and sleeps 5s (skips the poll entirely)
- `recover.sh` removes the marker at the very end (on both success and failure via `trap cleanup_inhibit EXIT`)
- This guarantees the watcher never counts failures during the intentional container-down window of a post-wake force-recreate

The marker file path is hardcoded to `/tmp/nullexit-warp-inhibit.marker` in both `toggle.sh` and `recover.sh`.

15.2.8 Incident Post-Mortem: Concurrency Collision between Nuclear Teardown and Post-Wake (July 10, 2026)
Symptom: When the background WARP Watcher triggers a nuclear shutdown (due to a real or simulated tunnel failure), the gateway is completely torn down. However, the network configuration changes made during the teardown trigger a `State:/Network/Global/IPv4` network change event. The LaunchAgent-based network watcher (`watcher.sh`) intercepts this event and concurrently spawns `recover.sh --post-wake` to heal/restore the gateway. This causes `docker compose down` (from nuclear teardown) and `docker compose up` (from post-wake) to run in parallel, resulting in container errors like `dependency failed to start: container warp exited (0)` and the gateway becoming wedged.

Root Cause: A lack of coordination/locking between the lifecycle control scripts. While `toggle.sh` used `/tmp/nullexit-toggle.lock` to prevent concurrent `toggle.sh` runs, `recover.sh` (both nuclear and `--post-wake` modes) did not utilize or check any lock files.

Fix (July 10, 2026): Implemented a shared lifecycle lock file (`/tmp/nullexit-toggle.lock`) across both scripts:
- **`recover.sh` (all modes)**: Checks `/tmp/nullexit-toggle.lock` at startup. If the lock is held by another running lifecycle script (`toggle.sh` or `recover.sh`):
  - In `--post-wake` mode, it logs a skip message to `output.log` and exits cleanly (`exit 0`). This safely prevents the auto-recovery agent from recreating containers during a teardown.
  - In nuclear recovery mode, it exits with an error.
- **`recover.sh` Lock Cleanup**: Cleans up the lock file upon completion using `trap cleanup_recover EXIT` (which checks if the lock file belongs to the current PID).
- **`toggle.sh` update**: The lock check in `toggle.sh` was updated to look for both `toggle.sh` and `recover.sh` in the running processes to enforce proper exclusion.

15.2.9 Incident Post-Mortem: Infinite Auto-Recovery Loop after WARP Watcher Nuclear Shutdown (July 10, 2026)
Symptom: When the background WARP Watcher triggered a nuclear shutdown (due to a tunnel outage), it successfully executed `recover.sh` (nuclear). However, immediately after the teardown completed, the system started spawning `recover.sh --post-wake` and rebuilding the containers again. This created an infinite loop of tearing down and auto-restarting the gateway, keeping the host's internet route permanently wedged.

Root Cause: An active-state marker file leak. The LaunchAgent-based network watcher (`watcher.sh`) only fires `recover.sh --post-wake` when `/tmp/nullexit-gateway-active.marker` is present on disk. While `toggle.sh` (STOP path) correctly deleted this marker, the nuclear `recover.sh` script did not. When the WARP Watcher ran `recover.sh` (nuclear), the containers were stopped but the active-state marker was left on disk. The network changes from the teardown then triggered `watcher.sh`, which saw the marker, believed the gateway was still supposed to be active, and triggered `--post-wake` to start it back up.

Fix (July 10, 2026): Modified the start of `recover.sh` to explicitly delete `/tmp/nullexit-gateway-active.marker` if `POST_WAKE` is `false` (nuclear mode). This ensures that any subsequent network configuration changes made during the teardown are correctly ignored by `watcher.sh` because the marker file is gone.

15.3 MTU, Fragmentation & Throughput

15.3.1 Network Bufferbloat & MTU Optimizations (Double-VPN UDP Crash)
Symptom: When running aggressive UDP bandwidth tests (such as macOS `networkQuality`, which uses QUIC/HTTP3), the `nullexit` tunnel completely crashes or exhibits severe lag (6,000ms+ bufferbloat).

Root Cause: The bug is a cascading MTU mismatch. Tailscale generates a UDP packet. Cloudflare WARP (also WireGuard) receives the packet, wraps it in another layer of encryption, pushing the packet size beyond the standard 1500 byte limit. Unlike TCP, UDP packets cannot be resized dynamically via MSS negotiation, resulting in massive IP packet fragmentation. The Apple Virtualization framework (`vz`), which Colima uses for bridged networking, silently drops heavily fragmented UDP packets under high load. This blackholes the entire UDP upload stream, instantly dropping the connection.

Fix (Partial): To prevent MTU fragmentation for standard web traffic, **TCP MSS Clamping** is strictly enforced via the macOS `pf` firewall. By adding `scrub out all max-mss 1160` to `scripts/pf.conf`, macOS unilaterally shrinks all outgoing TCP headers to fit comfortably inside the double-encrypted tunnel, bypassing Path MTU Discovery blackholes. `TS_DEBUG_MTU=1200` was also added to `docker-compose.yml` to minimize internal fragmentation. *(Superseded see 15.3.2: 1160 was later found still too large for double-tunneled traffic and tightened to 1120. That fix originally only reached the container's `post-rules.txt`; `pf.conf`'s host-side clamp stayed stale at 1160 until it was caught and synced to read `GATEWAY_MSS` dynamically, same as `post-rules.txt`, so the host and container can no longer drift apart.)*

Unresolved: The only way to solve the UDP crash bug natively is to artificially cap the maximum tunnel bandwidth using SQM (`fq_codel`). Because a static bandwidth cap would artificially throttle high-speed networks, `nullexit` leaves the bandwidth uncapped, optimizing perfectly for TCP while accepting UDP fragmentation under extreme edge-case load tests.
Why not Dynamic SQM (Autorate)? Implementing an EMA/PID-controlled autorate daemon (like `sqm-autorate`) requires a constant 100ms polling loop to measure the kernel packet queue, which severely degrades laptop battery life by preventing deep C-state CPU idling. Furthermore, macOS's native firewall (`dummynet`) only supports `fq_codel` and does not support the newer Linux `CAKE` algorithm, which is the only algorithm that offers kernel-native, event-driven `autorate-ingress` (zero polling penalty). Therefore, dynamic SQM is computationally hostile to battery-powered macOS hosts.

15.3.2 Double-Tunneling MSS Clamping (Stalling Bug)
Problem: Traffic double-wrapped through Tailscale (WireGuard) and WARP (WireGuard) suffered 120 bytes of MTU overhead (60 + 60). The default MSS clamp of 1180 was still slightly too large, causing large packets to fragment or drop, which resulted in mysterious web page stalls on strict-MSS endpoints.

Fix: Lowered the TCP MSS clamp in `post-rules.txt` to `1120` to guarantee double-tunneled payloads fit safely within standard internet MTUs.

15.3.3 Wi-Fi Edge Packet Loss via QUIC (UDP) Fragmentation
Problem: Applications that heavily utilize HTTP/3 (QUIC) over UDPsuch as Facebook, Instagram, and Google servicesexperience massive stuttering and stalled image loading when the client device is far away from the router (weak Wi-Fi signal). The issue did not occur when close to the router.

Root Cause: The `nullexit` gateway uses a double-encrypted tunnel (Tailscale + WARP). To prevent packets from fragmenting when entering these tunnels, a firewall rule in `post-rules.txt` uses `--set-mss 1120` to safely clamp TCP packets. However, because QUIC runs entirely over UDP, it bypasses the TCP MSS handshake completely. QUIC blasts maximum-MTU (1280 byte) UDP packets. The inner `tailscale0` interface (MTU 1200) forces these large UDP packets to fragment into two pieces. When the client is on the edge of Wi-Fi range (where 5-10% packet loss is common), losing *either* UDP fragment destroys the *entire* image frame. This exponential amplification of packet loss stalled QUIC streams.

Fix & Trade-offs: Appended an `iptables` rule to `post-rules.txt` that explicitly blocks `UDP --dport 443` (QUIC). When modern web apps detect QUIC is blocked, they instantly fall back to HTTP/2 (TCP 443). Because TCP is routed through the MSS clamp, the packets are resized *before* they leave, entirely eliminating IP fragmentation and restoring high-speed loading over weak Wi-Fi.
Consequences: This adds a minor latency penalty (TCP 3-way handshake) compared to QUIC's zero-RTT connection. It may also force WebRTC video calls (Google Meet/Discord) to fall back to TCP TURN servers, slightly inflating real-time video latency. However, in a constrained double-tunnel mesh, the absolute stability gained from eliminating fragmentation vastly outweighs these trade-offs.

15.3.4 Cellular Networks & PMTUD Blackholes (The 1280 Byte Limit)
Experiment: Conducted a live packet sweep from the PC-1 host to Phone-1 connected via a cellular network + Tailscale DERP relay using `ping -D -s <size>`.

Result:
- Packets with a payload of `1252` bytes (+ 28 bytes IP/ICMP headers = **1280 bytes total**) succeeded perfectly with ~60ms latency.
- Packets with a payload of `1253` bytes (**1281 bytes total**) and above resulted in a silent timeout.

Analysis: The cellular network (or the Tailscale DERP relay) enforces a strict MTU of 1280 bytes. Critically, it acts as a "PMTUD Blackhole" it silently drops packets larger than 1280 bytes instead of returning an ICMP "Fragmentation Needed" warning. If the exit node tries to send a standard 1500-byte internet packet to the phone over this link, it vanishes, causing web pages to stall permanently.

Conclusion: This empirically validates the absolute necessity of the TCP MSS clamping rule in `post-rules.txt`. By artificially clamping the MSS to `1120` (the value 15.3.2 above found necessary `1180` measured too large), we ensure the TCP payload + headers + double-WireGuard overhead never exceeds the strict 1280-byte ceiling of the cellular mesh link.

15.3.5 Low Gateway Throughput (DERP Relay & MTU Fragmentation) Fixed
Problem: The host Mac's internet throughput through the gateway dropped significantly (e.g., from 34Mbps raw to 2.1Mbps via gateway). This was caused by a combination of two bottlenecks:

1. Tailscale DERP Relay Bottleneck:
Because the host Mac and the Docker container operate behind the same public IP, standard hole-punching for Tailscale P2P failed, forcing traffic through Tailscale's NYC DERP relay. Normally, `routing-fix.sh` explicitly drops container Tailscale UDP packets destined for local subnets (when `TAILSCALE_ALLOW_LAN_P2P=false` in `.env` or auto-detected). We nuke these P2P connections on purpose to prevent the severe macOS `gvproxy` SNAT endpoint poisoning issue we faced earlier (see 15.7 SNAT Endpoint Poisoning). However, this strict policy unintentionally blocked direct communication between the container and the Mac host itself via the Docker bridge network.

Fix: Tailscale's control plane registers the Mac host using its physical IP (e.g., `172.17.52.223`), not the Docker bridge IP. If this physical IP falls within the dropped local subnets (like `172.16.0.0/12`), the P2P handshake is dropped. To solve this natively, `scripts/common.sh` now features a `write_host_ips()` function that actively grabs all physical IPv4 addresses on the Mac and writes them to `.host_ips`. `routing-fix.sh` dynamically reads this file every 30 seconds and injects priority `RETURN` rules for those exact physical IPs. This guarantees direct container-to-host P2P over the local bridge (bypassing DERP) while continuing to block external LAN devices (like phones) to prevent the 15.7 SNAT poisoning bug.

2. MTU Double-Encapsulation Fragmentation:
The host's Tailscale interface (`utun5`) defaulted to an MTU of 1280. The WARP tunnel inside the container also uses an MTU of 1280. When a full 1280-byte Tailscale packet from the host entered the container and was encapsulated by WARP (adding ~60 bytes of overhead), the resulting packet exceeded 1280 bytes, leading to severe fragmentation and performance degradation.

Fix: In `toggle.sh`'s `setup_exit_node_routing` function, the host's Tailscale `utun` interface MTU is explicitly lowered to `1200` (configurable via `HOST_MTU` in `.env`). This ensures that host-originated packets can comfortably fit inside the container's WARP tunnel encapsulation without fragmenting.

> **Design Note Is direct hostcontainer P2P a Docker bug, or are we "exploiting a VM/host relationship"?**
> Neither. It is expected, well-defined behavior, and the RETURN-rule carve-out is a *defensive de-restriction*, not an exploit. The reasoning:
> - **On native Linux**, the container and host share one kernel; the Docker bridge is a first-class local interface and hostcontainer is *trivially* direct. Nobody calls that an exploit. Our RETURN rules simply reproduce that same local reachability on macOS.
> - **The RETURN rules add nothing new; they *stop blocking* something legitimate.** `routing-fix.sh` broadly DROPs containerLAN Tailscale UDP to prevent the real 15.7 `gvproxy` SNAT endpoint-poisoning bug caused by *external* LAN peers (e.g. a phone on another AP). Tailscale's control plane registers the Mac host by its *physical* LAN IP (e.g. `172.17.52.223`), which falls inside those DROPped ranges so the host, the one peer that is literally the same physical machine, became collateral damage. The `.host_ips` RETURN rules are the minimal, precise exception that re-permits exactly that same-machine path and nothing else.
> - **The genuinely VM-specific wart is `gvproxy` SNAT poisoning (15.7), and we route *around* it, not through a hole in it.** Keeping hostcontainer traffic on the Docker bridge (direct) avoids letting it get NATed through the userspace `gvproxy` endpoint rewriter that confuses Tailscale's peer-endpoint learning. That is the opposite of exploiting the VM boundary it is respecting it.
> - **What genuinely *doesn't* work on a VM host is UDP hole-punching to arbitrary *external* peers** (see 15.13). That is a real limitation of userspace VM networking, and it is why remote peers fall back to DERP unless bridged mode is enabled on a trusted LAN. Direct hostcontainer P2P is the one case that works regardless because both endpoints live on the same physical machine, no hole-punching is needed at all. It was verified `direct 172.17.52.223` even in plain user-mode networking (no `--network-address`).
>
> In short: the DROP is the deliberate policy, the RETURN is a surgical exception for same-machine traffic, and the whole thing is standard bridge networking not a Docker bug and not an exploit.

15.3.6 The "~4Mbps Ceiling" Was a Benchmark Artifact, Not a Tunnel Bottleneck (July 18, 2026)
Symptom: Repeated throughput tests against `speedtest.tele2.net` from both the host (full double-tunnel) and the `warp` container (isolated via its own SOCKS5 proxy) consistently landed around 3.3-4.7 Mbps, regardless of which leg was tested. This looked like a hard nullexit-imposed ceiling and was the leading suspect for the Instagram/general-slowness symptoms in 15.3.3-15.3.4.

Investigation: Before accepting "the tunnel caps at ~4Mbps," the same target was measured **raw** gateway fully stopped, plain ISP egress, no tunnel at all and it *still* came back at ~3.2 Mbps. A second raw test against a fast, nearby CDN (Cloudflare) returned ~32.2 Mbps on the same physical connection. `speedtest.tele2.net` was the bottleneck, not nullexit; it's a slow/distant target that happened to be the first thing tried.

Confirmation: Re-measured through the tunnel against a fast CDN that (unlike Cloudflare's own speed-test endpoint, which 403s WARP's shared/CGNAT egress IPs as abuse-prevention) doesn't flag WARP IPs Hetzner's public speed-test files. Result: **15.8-26.4 Mbps** tunneled, in the same ballpark as the 32.2 Mbps raw baseline (roughly 20-50% overhead from the double WireGuard encapsulation, normal for this architecture). CPU and memory on the Colima VM were also checked during sustained transfers (`colima ssh -- top -bn1` / `free -m`) and ruled out as a bottleneck load average stayed near 0.00-0.07 and swap held flat around 250MB throughout, on both a run that completed cleanly and one that reset mid-transfer.

Fix: `scripts/sweep.sh`'s throughput check (9, check 7/7) now defaults to the Hetzner target instead of `speedtest.tele2.net`, so future sweeps report real tunnel performance instead of accidentally re-measuring a slow third-party server and misattributing it to nullexit.

Residual, unchanged: the occasional mid-transfer connection reset (curl exit 56, `Recv failure: Connection reset by peer`) observed on maybe 1-in-6 to 1-in-8 attempts is still unexplained by CPU/memory and remains consistent with the gvproxy/UDP-fragmentation-under-load issue in 15.3.1 a real, separate, and still-open question from the throughput-ceiling one this section resolves.

15.3.7 Field Test: Dense Cellular Congestion on a Mesh Device (July 19, 2026)
Setup: An Android mesh device on cellular data, physically present in a very dense public crowd, exit-noded through `nullexit-gateway` for the whole session. Goal: see whether heavy local RF/tower congestion degrades the tunnel differently than normal cellular use.

Method: Confirmed the device was actively routing live (not just tailnet-idle) via AdGuard's query log showing fresh DNS lookups from it seconds before the test. Then ran a large TSMP burst (`tailscale ping --until-direct=false`, the same mechanism as sweep.sh check 6/7 but a much larger sample) to get a real loss/jitter signal instead of a single ping.

Result: Every probe got a pong back 0% loss. The path was DERP-relayed for the entire burst, never direct, which is expected on cellular (carrier-grade NAT blocks Tailscale's UDP hole-punching regardless of crowd density, so this isn't specific to a busy area). RTT stayed in a tight, low band for the large majority of probes, with a couple of outlier spikes consistent with intermittent RF contention/retransmission at the cell tower from the sheer device density, not a gateway-side problem. All 5 containers stayed healthy throughout; the WARP double-tunnel was confirmed live.

Conclusion: nullexit was not the bottleneck under real-world dense-cellular congestion zero packet loss is a strong result. The DERP-vs-direct baseline and the latency spikes both trace to the cellular/CGNAT layer, outside nullexit's control.

15.4 DNS Interception & Proxy

15.4.1 DNS Proxy: TCP Wire Format Mismatch (socat / dnsmasq)
Problem: When the Tailscale data plane is unavailable (DERP relay only), the script falls back to a local DNS proxy that forwards queries from host UDP:53 to Docker's TCP:5354 (AdGuard). The `socat` and `dnsmasq` approaches both failed.

- **socat** forwards raw bytes it does not prepend the 2-byte length prefix that DNS-over-TCP requires. AdGuard parses the stream incorrectly and returns garbage.
- **dnsmasq** tries UDP first to reach the upstream (`127.0.0.1#5354`), but Colima's SSH tunnel only forwards TCP. With a single upstream server, dnsmasq doesn't fall back to TCP even with `--timeout=1`.

Solution: Replaced both with a **Python DNS proxy** (`scripts/dns-proxy.py`, ~25 lines) that properly handles the DNS-over-TCP wire format: reads a UDP query from the host, prepends the 2-byte length prefix, sends it over TCP to AdGuard, strips the prefix from the response, and sends it back over UDP.

Architectural Note: Why is the DNS proxy in Python while the SOCKS5 proxy was moved to a compiled Go Docker container?
The SOCKS5 proxy handles heavy data streaming (e.g., video, downloads). Python introduces massive CPU/battery overhead for raw socket streaming, which is why we rely on the compiled Go implementation built directly into the `tailscaled` daemon via `TS_SOCKS5_SERVER=:1080`.
Conversely, the DNS proxy only handles intermittent UDP queries (a few bytes per minute) generated solely by the local host machine. The Python overhead for this is effectively 0.001 Watts. We deliberately kept `dns-proxy.py` in Python because it runs natively on the macOS/Linux hostrewriting it in Go would force users to install a Go compiler (`brew install go`) on their host machine for zero noticeable battery gain. (Note: If this architecture is ever adapted to serve as a public DNS resolver for thousands of users or for massive automated web-scraping clusters, `dns-proxy.py` must be rewritten in Go to survive the concurrent packet flood).

15.4.2 Tailscale MagicDNS Bypassing the AdGuard PREROUTING Intercept
Context: The `routing-fix.sh` script applies an iptables NAT rule (`PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT`) to intercept all incoming DNS queries from connected Tailscale peers and force them into the AdGuard container on port 5335.
Problem: When remote clients (like Android or iOS) have "Use Tailscale DNS settings" turned on, they query Tailscale's MagicDNS IP (`100.100.100.100`). This packet enters the exit node, but the `tailscaled` daemon running *inside* the gateway container intercepts the `100.100.100.100` packet internally. `tailscaled` then generates a brand new, local DNS request to the upstream DNS server to resolve the query. Because this new request is generated *locally* by the daemon (not arriving externally via the `tailscale0` interface), it bypasses the `PREROUTING` chain completely. The request glides straight out the WARP tunnel to Cloudflare/Google, entirely bypassing the AdGuard sinkhole.
Fix: The only way to fix this blindspot is to force `tailscaled` to use AdGuard as its upstream. In the Tailscale Admin Console, the gateway's Tailscale IPv4 address (e.g., `100.x.x.x`) must be added as the sole Custom Global Nameserver, and "Override local DNS" must be enabled. Additionally, to block Android devices from bypassing this via Private DNS (DNS-over-TLS), outbound Port 853 traffic is explicitly REJECTED in both `routing-fix.sh` and `post-rules.txt`.

15.4.3 Seamless Wi-Fi Roaming & The macOS DNS Wipeout Loophole
Problem: The Tailscale, WireGuard, and Colima NAT stacks are natively designed for connectionless roaming and will seamlessly survive when the macOS host switches Wi-Fi networks (e.g., roaming from home Wi-Fi to a cellular hotspot). However, the internet completely breaks upon network transition. This occurs because macOS intentionally wipes out all custom DNS settings (`networksetup -setdnsservers`) and reverts to the new network's default DHCP nameserver whenever the active Wi-Fi BSSID changes. Since the Mac is locked to the exit node, its DNS requests are swallowed by WARP and dumped onto the public internet, which cannot route private DHCP IPs. This results in a total DNS failure.

Solution: Implementing a silent background "DNS Watcher" daemon (`nullexit-dns-watcher`) in `toggle.sh`. When the gateway starts, it spawns a background polling loop that checks the active Wi-Fi DNS every 30 seconds. If macOS resets the DNS during a network change, the watcher instantly detects the drift and forcefully re-injects `100.100.21.8` via `networksetup`. When the gateway stops, the watcher process is cleanly killed. This allows true, seamless roaming across physical networks without ever dropping the gateway state.

15.4.4 docker compose exec `\r` Injection
Fix: All `docker compose exec` output piped through `tr -d '\r'`.

15.4.5 Stderr Invisibly Swallowed
Symptom: Failures were silent due to stderr redirected to `output.log`.

Fix: Tailscale wait loop now shows output. Critical commands show errors inline.

15.5 Roaming, Sleep/Wake & Auto-Recovery

15.5.1 Gateway Breakage on Lid Close / Wi-Fi Roam Post-Wake + Post-Roam Auto-Recovery
Context / Symptom. Closing the lid (sleep/wake) OR losing + reconnecting Wi-Fi (e.g. in an elevator, caf, or roaming between APs) leaves the gateway in a partly-dead state. Symptoms range from a ~30s window of dead DNS to a permanent outage that only full `recover.sh` or `toggle.sh` fixes. The root cause is that nullexit has **no daemon watching macOS sleep/wake or network-state-change events** the existing DNS

Watcher inside `toggle.sh` only runs in-process and is suspended along with the rest of the user shell on lid close, so the gateway cannot self-heal after either event.

Diagnosis. Two distinct failure modes share the same root gap.

* **(A) Lid close sleep wake.** `caffeinate -i` (used by `toggle.sh`) blocks *idle* sleep but **does not block forced sleep from lid close**. Closing the lid hard-suspends Colima VM + every container + every Tailscale-quit-shells-except-tailscaled. On wake the VM clock needs ~5-30s to resync via NTP, during which TLS cert validation fails: Cloudflare WARP `162.159.192.1:2408` rejects the handshake gluetun's healthcheck (`interval: 2s, retries: 15`) eventually flips unhealthy Docker `unless-stopped` restarts the `warp` container (~30s gap). `mDNSResponder`'s in-memory cache still holds pre-sleep entries (stale A records for tailnet hosts) so the first dozen DNS queries after wake time out. `tailscaled` on the host often keeps its stale DERP relay preference and routes packets through a dead relay for another 30-60s.
* **(B) Wi-Fi roam / loss in elevators, captive portals.** WARP is a UDP tunnel to Cloudflare's anycast edge. Carrier NATs and captive-portal networks silently invalidate the UDP binding on roam and on hotspot-bound re-association. macOS `networksetup` *should* preserve DNS settings on the same Wi-Fi service across a roam, but most captive-portal networks DHCP-replace DNS to the captive-portal resolver **during the captive-portal dance itself**, before the user has authenticated so even DNS hijack survives a clean roam and quietly breaks on captive-portal handoff.

Resolution. A single **launchd LaunchAgent** (`com.nullexit.wake-recovery`) runs a long-running shell daemon (`scripts/watcher.sh`) that listens for both event sources and, on each fire, executes `bash recover.sh --post-wake` a *lighter-touch* recovery mode that's the opposite of the existing `--nuclear` semantics: it keeps the gateway live while still refreshing every stale subsystem.

Deep dive: Apple Power Management Primer (for the unfamiliar).
macOS exposes several relevant surfaces; the right primitive depends on what you actually need to detect. None of these are obvious and none of them have a standard splash screen in the docs.

* **`caffeinate <flags>`** creates I/O Kit power assertions via `IOPMAssertionCreateWithName`. It does NOT block forced sleep (lid close, low-battery shutdown, sleep timer firing on a per-set Energy Settings policy). Flags:
  * `-i` PreventIdleSleep (the only one `toggle.sh` uses; protects against the OS going idle while the user is active but a clamshell close still hard-puts it to sleep)
  * `-s` PreventSystemSleep (only honored on AC power; the user may be on battery)
  * `-u` DeclareUserActive (resets the idle timer every 30s by default; equivalent to keeping the cursor moving)
  * `-d` PreventDisplaySleep
  * `-m` PreventDiskIdleSleep
  * There is NO `-l` (lid-block) flag in any modern macOS. Lid close is a kernel-firmware-level event that ignores user-space assertions.
  * To prove any of this on live hardware: `pmset -g assertions` lists every active assertion with type / process / reason / timeout.
* **`pmset`** is the read-side of the same system:
  * `pmset -g log | grep -E 'Sleep|Wake|DarkWake'` shows the recent timeline.
  * `pmset -g history` is the per-day summary.
  * `pmset -g assertions` is the live assertion list (matches `-i -s -d -u` to processes).
* **`launchd` does NOT have a native "on-wake" hook.** Not in any version of macOS as of Sonoma. The canonical recipe is `StartCalendarInterval` with a sub-minute cadence and let launchd queue missed intervals for replay-on-wake, or use a third-party daemon (Sleepwatcher). For our use case a long-running shell daemon listening to the unified log is more responsive than a polling agent.
* **`com.apple.system.powermanagement.*` Darwin notifications** (`willSleep`, `didWake`, `didDim`, `didUndim`) are the C-level signal. From a shell, they are best reached through the unified log with `log stream --predicate` see the watcher implementation.
* **`scutil n.watch`** is the canonical CLI for live-following network-state changes. Pairs with `n.add State:/Network/Global/IPv4` to get a single tap that fires on **any** IPv4 link/route/DNS/address change across **all** interfaces. This is what `scripts/watcher.sh`'s network-change listener uses.
* **`com.apple.networkChange` Darwin notification** is the C-level equivalent but has no shell-friendly listener; use `scutil` instead.
* **`SCNetworkReachability`** is the Objective-C API used by SOCKS5/HTTP clients to detect route changes per target. Never a fit for shell scripts.

The Fix (component-by-component).

1. **`recover.sh --post-wake`** (new flag, **non-destructive** by design). Adding `--post-wake` to the existing nuclear recovery script inverts the semantics. The default mode still tears down Tailscale, resets DNS to empty, stops caffeinate, runs `docker compose down`, power-cycles Wi-Fi. The new mode does:
   * Skip Tailscale disconnect (we WANT it to keep the mesh connection)
   * **Re-hijack** DNS to the gateway Tailscale IP read from `.gateway_ip` (instead of resetting to empty) applies to both `ACTIVE_SERVICE` and `EN0_SERVICE` because the system resolver is scoped to en0
   * **Refresh** the exit-node preference with `tailscale up --reset --ssh=true --accept-dns=false --exit-node=\"$TS_IP\" --exit-node-allow-lan-access=true` (the exact flags `toggle.sh` START uses). This re-asserts DERP mapping without dropping the mesh
   * Inspect `docker inspect --format '{{.State.Health.Status}}' nullexit-warp-1` and **only** `docker compose up -d --force-recreate warp` if gluetun is unhealthy this single targeted recreate nudges the UDP NAT binding back to Cloudflare without disturbing Tailscale or AdGuard
   * Run the sharingd-reset step (existing fix from `toggle.sh` START path; see 15.11) so AirDrop doesn't freeze on the new IP lease
   * Verify the gateway with `dig +tcp @127.0.0.1 -p 5354 google.com` (AdGuard via TCP through Colima's SSH tunnel) and a 4-second `tailscale ping` to the gateway Tailscale IP, instead of the default mode's direct-to-1.1.1.1 check

```
  * Crucially: **skip `sudo -v`** because the LaunchAgent has no TTY to prompt on; all privileged calls use `sudo -n` and lean on `/etc/sudoers.d/nullexit` NOPASSWD entries (`networksetup`, `dnsmasq`, `dscacheutil`, `killall`, `pkill`).
2. **`toggle.sh`** writes and clears `/tmp/nullexit-gateway-active.marker` (an ISO-8601 UTC timestamp) at the bottom of the START path and the top of the STOP path / `cleanup_handler`. Two new helpers: `write_gateway_active_marker` and `clear_gateway_active_marker`. The watcher uses this file as its **only signal** that the gateway is live: if `toggle.sh` was stopped (or never started), the watcher no-ops. This prevents waking-up-only-on-counterfactual events from churning DNS.
3. **`scripts/watcher.sh`** (long-running daemon, sibling-style source file). Two backgrounded pipelines:
  * **Wake listener**: `log stream --predicate 'composedMessage CONTAINS[c] "didWake" OR composedMessage CONTAINS[c] "Wake from Sleep" OR composedMessage CONTAINS[c] "Waking from" OR composedMessage CONTAINS[c] "DarkWake"' --style compact --info` piped through `while read; do run_recover "WAKE: ..."; done`. `[c]` makes the predicate case-insensitive (the unified log writes phrases in mixed case across versions).
  * **Network listener**: `(echo "n.add State:/Network/Global/IPv4"; echo "n.watch"; while true; do sleep 86400; done) | scutil 2>/dev/null` piped through a `case` match on `n.state*` / `SCEventUpdate*`.
  * **`run_recover`** checks the marker, debounces via `/tmp/nullexit-watcher.last-recovery` (default 10s, configurable via `NULLEXIT_DEBOUNCE_SECONDS`), and shells out to `bash recover.sh --post-wake`. The 10s debounce prevents a single wake event (which typically emits 2-3 log lines) from triggering multiple recoveries.
  * `trap ... TERM INT`  `kill 0` cleanly tears down the process group on launchctl `bootout` so the `sleep 86400` inside the scutil pipe also dies.
  * `wait` blocks indefinitely while both pipelines feed it.
4. **`launchd/com.nullexit.wake-recovery.plist`** (LaunchAgent in the user domain  runs as the console user at every login without needing sudo).
  * `Label: com.nullexit.wake-recovery`
  * `ProgramArguments: ["/bin/bash", "<absolute-path-to-your-nullexit-install>/scripts/watcher.sh"]`
  * `RunAtLoad: true`  starts immediately on `launchctl load` (no need to log out and back in)
  * `KeepAlive: { SuccessfulExit: false, Crashed: false }`  **don't** restart on clean SIGTERM exit (so `launchctl bootout` doesn't get stuck in a relaunch loop). Also **don't** restart on hard crash: we observed that macOS Sonoma+'s sleep/wake suspend-resume path accumulates orphan watcher.sh PIDs because SIGCONT does not reliably clean up the prior instance's listener grandchildren, and a `KeepAlive.Crashed=true`-driven relaunch actively races with the still-live prior process. The script enforces single-instance on its own via a `flock`-style PID-file lock, so a relaunch-on-crash would be skipped by the lock and produce 0 listener coverage until the live instance happened to die. Trade-off: a real crash leaves the user without auto-recovery until they run `launchctl load -w` manually  acceptable because (a) gateway still works without auto-recovery, (b) a relaunch wouldn't fix whatever caused the crash, and (c) the gateway-recovery itself is observably broken (no DNS, no exit-node) if it does go down.
  * `ProcessType: Background`  hint to App Nap not to suspend us.
  * `StandardOutPath/StandardErrorPath: /tmp/nullexit-watcher.{out,err}.log`  separates launchd-level diagnostics from the script's own `/tmp/nullexit-watcher.log` (which `exec >>` redirects).

**Install (one-time).** The plist lives in the repo at **`launchd/com.nullexit.wake-recovery.plist`** for version control. The installed (live) copy must land in `~/Library/LaunchAgents/` so launchd picks it up on the user session.

```bash
# Copy the plist to the user-launchd location (run from repo root)
cp ./launchd/com.nullexit.wake-recovery.plist \
   ~/Library/LaunchAgents/

# Substitute __NULLEXIT_HOME__ with your local install absolute path.
# The plist uses WorkingDirectory + a relative program arg, so this
# single substitution is the only place it references your machine.
sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" \
    ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Validate the XML
plutil -lint ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Load + persist this user session + every future login
launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Confirm it actually started
launchctl list | grep nullexit
```

To **stop** the watcher (without uninstalling):

```bash
launchctl bootout gui/$UID/com.nullexit.wake-recovery
# or:
launchctl unload -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```

To **uninstall** completely:

```bash
launchctl bootout gui/$UID/com.nullexit.wake-recovery 2>/dev/null || true
rm ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```
```

Why this fix was preferred over alternatives.

* **Polling with `StartCalendarInterval`** would have given ~30-60s detection latency per wake. The unified-log stream is sub-second.
* **Sleepwatcher** would have covered wake events but not network changes. We'd still need a separate `scutil` listener, and introducing a third-party dependency for what amounts to ~150 lines of shell is unjustified.
* **Forcing `caffeinate -l`** doesn't exist and the closest equivalent (`caffeinate -s`) only works on AC power. The whole point of the gateway is to stay usable on battery while traveling lid close must NOT silently kill the gateway.
* **A kernel extension** is impossible (no entitlement) and out of scope.

Reproduction recipe. Test both events in isolation to confirm the fix actually closes the gap.

* **(A) Lid-only repro**: with the gateway up, run in a terminal: `pmset displaysleepnow && sleep 30 && pmset wake` (without physically closing the lid). Watch the wake wake-fires the watcher post-wake routine runs login should still work in <5s after wake. Compare to baseline pre-fix: pre-fix the gateway would be dead for 30-90s.
* **(B) Roam-only repro**: with the gateway up: `networksetup -setairportpower off && sleep 30 && networksetup -setairportpower on` (or simply walk out of Wi-Fi range and back). The captive-portal WILL repave DHCP DNS for a few seconds; the watcher fires on the second `State:/Network/Global/IPv4` change (after the Wi-Fi re-associates) and the post-wake routine re-hijacks DNS within 2-3s.

Operative files.

* `recover.sh` added arg parsing; wrapped destructive sections behind `POST_WAKE`; added re-hijack DNS / refresh exit-node / force-recreate-if-unhealthy path.
* `toggle.sh` added `write_gateway_active_marker` / `clear_gateway_active_marker` called at the END of START and TOP of STOP / cleanup_handler.
* `scripts/watcher.sh` new long-running daemon.
* `launchd/com.nullexit.wake-recovery.plist` launchd LaunchAgent descriptor.
* `~/Library/LaunchAgents/com.nullexit.wake-recovery.plist` the installed copy (one-time copy).

Honest assessment (read before testing). This fix has two asymmetric halves:

* **Network / roam / captive-portal path RELIABLE.** `scutil n.watch` on `State:/Network/Global/IPv{4,6}` catches every Wi-Fi roam, captive-portal DHCP-rebind, hotspot handoff, and VPN change with zero missed events. To validate, drop Wi-Fi for 30 s and re-add it; `/tmp/nullexit-watcher.log` will record a `NET:` trigger and `recover.sh --post-wake` will run within ~10 s of the network coming back.
* **Lid close / wake path BEST-EFFORT on macOS Sonoma+.** Apple *suspends* LaunchAgents process-state during forced sleep and *resumes* them on wake; it does NOT re-launch them every wake. As a result: (a) `log stream` is a live subscription events emitted during the agent's suppressed window are DROPPED, not buffered, so the wake event that prompted the resume is invisible to `log stream` on resume; (b) the "fire on startup" guard runs once on `launchctl load -w` and after a `KeepAlive` crash-recovery, NOT on every successive wake; (c) `subsystem == "com.apple.powermanagement"` will catch FUTURE emissions (display dim/restore, manual sleep/wake) but not the lid-closeopen wake directly.

In practice the lid-wake gap is short, because `toggle.sh`'s existing DNS Watcher already polls every 5 s and re-hijacks DNS (so DNS recovers within 5 s of any roam or wake), `tailscaled` self-heals via DERP fallback within 3060 s, and `gluetun` reconnects via its healthcheck retry within 30 s. The user-visible pain on lid open is ~30 s of wonky DNS, not a permanent outage.

Test the ROAM path FIRST. If ROAM works in your testing, the lid-wake gap is acceptable. If lid-wake handling is critical, the two clean follow-ups are:
1. **Sleepwatcher** (one canonical install): `brew install sleepwatcher`; drop `recover.sh --post-wake` into `~/.sleep` and `~/.wake` so Sleepwatcher invokes it directly on every wake without the launchd propagation problem.
2. **Move the watcher to a LaunchDaemon** (root) and use `IOPMSchedulePowerEvent` via a tiny C wrapper not portable to shell.

A truly reliable shell-only solution to "wake events from inside a LaunchAgent" does not exist on macOS Sonoma+.

Empirical lessons from deployment. Two findings came up while deploying this fix end-to-end; recording so the next operator doesn't lose an evening to each.

1. **Bash function-call-before-definition gotcha.** While introducing the gateway-active marker helpers, the defensive `clear_gateway_active_marker` call was inserted at the very top of `toggle.sh` (right after `set -e`), but the function definition lived near `PID_FILE="/tmp/nullexit-caffeinate.pid"` ~90 lines down. Bash parses top-to-bottom and resolves names at first invocation, so the call site exited with code 127 (`clear_gateway_active_marker: command not found`). The gateway stayed unreachable from the START path even though `bash -n` passed (the parser doesn't catch order-of-resolution errors). Rule: define functions BEFORE any block that calls them. Verify with `grep -n '^name()\s*{'` followed by `grep -n '^name\s*$'` the latter must appear on a line number AFTER the former.
2. **`networksetup -setairportpower off+on` is not a faithful roam reproduction.** Synthetic Wi-Fi radio-cycling (`sudo networksetup -setairportpower Wi-Fi off` for 30 s, then back on) is expected to produce a `NET:` trigger in `scutil n.watch` but produced ZERO during testing because `setairportpower` powers the radio at the driver level while leaving the network SERVICE in the `"up"` state from scutil's perspective, so no `n.state State:/Network/Global/IPv4` event is generated. A real roam (macOS briefly loses the AP and re-associates into a different one) DOES fire the event because the link actually changes.

   Better reproduction recipes in priority order:

```
1207    * Walk between two real SSIDs via `networksetup -switchtolocation` (or actual physical station movement) guaranteed to fire the scutil event.
1208    * Force a DHCP rebind on the same SSID with `sudo ipconfig set en0 DHCP` triggers `State:/Network/Global/IPv4` to update.
1209    * Trust the existing 30-second DNS Watcher inside `toggle.sh` for the DNS path: it re-hijacks DNS automatically. The post-wake watcher adds clear value mostly for tailscale exit-node re-assertion and warp gluetun force-recreate, both of which self-heal within ~30 s anyway.
1210
1211 #### 15.5.2 July 8, 2026: Network Watcher Event Mismatch, Sudo-in-Background Crash, and Loop-Prevention
1212 **Symptom:**
1213 1. Switching Wi-Fi or waking the Mac from sleep did not trigger a post-wake recovery (`recover.sh --post-wake`) automatically.
1214 2. The network connection would eventually get completely sinkholed/blocked. If the user tried to restart or wait, it did not self-heal.
1215 3. During the recovery attempt, checking the public IP on the host would return the unencrypted campus/ISP IP (starting with `132.`) instead of the encrypted tunnel IP.
1216
1217 **Root Cause:**
1218 - **Bug 1 (Network Watcher Mismatch):** In `scripts/watcher.sh`, the network change listener watched `State:/Network/Global/IPv4` changes via `scutil n.watch` but matched them against `*n.state*|*SCEventUpdate*`. On macOS, the actual output is `changedKey [0] = State:/Network/Global/IPv4`. Because the pattern did not match, all network switches were silently ignored.
1219 - **Bug 2 (Sudo Caching Background Crash):** When the background WARP Watcher eventually detected the broken tunnel (after 6 failures/30s), it triggered the default nuclear `recover.sh` (with `POST_WAKE=false`). Because there was no active TTY in the background context, the `sudo -v` caching check aborted instantly with a terminal required error. Due to `set -e`, `recover.sh` died immediately before resetting DNS, disabling the packet-filter killswitch, or stopping containers, leaving the host network completely sinkholed.
1220 - **Bug 3 (Post-Wake Infinite Loop):** Once the watcher patterns were fixed, `recover.sh --post-wake` triggered successfully on network changes. However, because the script deleted the default routing entries at the start, spent ~15 seconds rebuilding 88 DERP relay bypass routes, and re-added the default routes at the end, this execution time exceeded the watcher's 10-second debounce threshold. Furthermore, the watcher recorded the *start* timestamp in the debounce file. Consequently, the route changes made by the script itself triggered the watcher again immediately after completion, trapping the system in an infinite loop of recovery runs. During this loop, the VPN routes were constantly deleted, resulting in the host leaking traffic through the real ISP interface (an IP starting with `132.`) for 15s out of every cycle.
1221
1222 **Resolution:**
1223 - **Fixed Watcher Matching:** Updated `scripts/watcher.sh` to match `*changedKey*` and `*State:/Network/Global/IPv*` to correctly catch network changes.
1224 - **Bypassed Background Sudo:** Added `--auto`/`--non-interactive` flags to `recover.sh` and added a TTY check `[ -t 0 ]` to skip interactive `sudo -v` authentication when running headlessly in the background. Updated the WARP Watcher call in `toggle.sh` to pass `--auto`.
1225 - **Prevented Infinite Loops:** Modified `scripts/watcher.sh` to write the **completion** timestamp of `recover.sh` to `DEBOUNCE_FILE` (instead of only the start timestamp). This ensures all buffered network config events generated during route reconstruction are evaluated against the end-of-run timestamp and successfully debounced, breaking the infinite loop.
1226 - **Centralized & Timestamped Logs:** Redirected `watcher.sh` stdout/stderr from `/tmp/nullexit-watcher.log` to the repository's main `output.log` and updated `scripts/common.sh`'s `step`, `ok`, `warn`, `fail`, and `die` helpers to prefix logs with a local `[YYYY-MM-DD HH:MM:SS]` timestamp.
1227
1228 #### 15.5.3 Incident Post-Mortem: Startup Pre-Flight Check Tailscale Race (July 10, 2026)
1229 **Symptom:** Immediately after startup via `toggle.sh`, the pre-flight connectivity check `[1/3] Gateway reachable via Tailscale` returned `FAIL`, forcing `toggle.sh` to skip the full exit-node activation and fall back to SOCKS5/local DNS proxy mode.
1230
1231 **Root Cause:** A timing race condition during startup. In `toggle.sh`, the host joins the mesh using `tailscale up --reset` (Phase A) right before performing the pre-flight check. Because `tailscale up` resets and reconfigures the host's `tailscaled` daemon, the local daemon must fetch the latest netmap asynchronously from the coordination servers. This network map propagation and route establishment takes a few seconds. Running `tailscale ping` immediately after `tailscale up` returns (with no delay) causes the check to fail because the local daemon has not yet discovered the gateway node `100.100.21.8`.
1232
1233 **Fix (July 10, 2026):** Upgraded pre-flight Check 1 in both `toggle.sh` and `scripts/toggle-linux.sh` to use a 5-attempt retry loop with a 1-second delay between checks. This allows up to 5 seconds for local tailnet routing paths to settle, ensuring the gateway is correctly reached and a pong is received before deciding whether to enable exit-node routing.
1234
1235 #### 15.5.4 Incident Post-Mortem: Pre-Flight Check False Negatives due to VPN Firewall STUN Blocks (July 10, 2026)
1236 **Symptom:** Immediately after a cold boot or full restart of the gateway, the pre-flight connectivity check `[1/3] Gateway reachable via Tailscale` failed, causing `toggle.sh` to skip the full exit-node configuration and fall back to SOCKS5/local DNS proxy mode. However, manually running `tailscale ping 100.100.21.8` immediately after startup succeeded in 1ms.
1237
1238 **Root Cause:** An asynchronous Tailscale handshake delay caused by firewall restrictions. Inside the container network namespace, the `warp` container (gluetun) implements a strict firewall that blocks all outbound UDP traffic to the public internet except to the designated WireGuard endpoint. Because of this, the `tailscale` container's STUN requests to public STUN servers are blocked, causing the local daemon to report `UDP is blocked, trying HTTPS`. Without UDP STUN capability, Tailscale is forced to fall back to a slower DERP-assisted handshake (relayed via HTTPS/TCP). This negotiation and direct-path punch-through takes
```

roughly 30 to 35 seconds to establish on cold boots (after a full `tailscale up --reset`). Since the pre-flight check only waited 15 seconds (15 attempts with 1-second sleeps), the check timed out and aborted before the path completed.

Fix (July 10, 2026): Increased the pre-flight ping retry limit from 15 to 45 attempts (45 seconds) in both `toggle.sh` and `scripts/toggle-linux.sh`. This provides sufficient time for the DERP-assisted handshake to complete on cold boots, while still passing immediately (in 1-2 seconds) on warm starts or once the path is established.

15.5.5 Incident Post-Mortem: Docker Image Build DNS Failures on Immediate Restart (July 10, 2026)

Symptom: When executing `toggle.sh --restart`, the script successfully stops the gateway but then immediately fails during startup during `docker compose build` with DNS resolution errors: `failed to do request: Head "https://registry-1.docker.io/...": dial tcp: lookup registry-1.docker.io on 192.168.5.1:53: no such host`.

Root Cause: A race condition between physical link negotiation and virtual machine startup. The STOP path of the restart command invokes `cleanup_network_state` which power-cycles the host's physical Wi-Fi interface (`en0`) to flush stale routing states. Immediately after, `toggle.sh` starts the gateway boot sequence. Because `toggle.sh` did not verify the status of the physical interface, it proceeded to boot Colima and build Docker containers while `en0` was still negotiating a DHCP lease. Consequently, the host (and by extension, the VM bridge) had no active internet gateway/DNS path, causing Docker's registry check to fail.

Fix (July 10, 2026):

1. Created a unified `wait_for_dhcp_settle` helper function in `scripts/common.sh` that polls `ipconfig getpacket` and `ifconfig` until a valid IP and router are assigned. To handle typical macOS Wi-Fi reassociation and DHCP negotiation latency (which commonly takes 15-20 seconds after a power-cycle), this loop was configured to poll up to 60 times (30 seconds total).
2. Injected a call to `wait_for_dhcp_settle` at the beginning of the `START` action in `toggle.sh` (right before checking Colima status) to block container builds until the host's physical network is online.
3. Refactored `recover.sh` to reuse the unified `wait_for_dhcp_settle` function.

15.5.6 Watcher.sh Suppressed Exit Code Bug (Fixed)

Observation: The `scripts/watcher.sh` script is a long-running daemon that invokes `recover.sh --post-wake` when it detects system wake or network roam events. Previously, it contained the following bug:

```bash
  bash "$RECOVER" --post-wake
  local exit_code=$?
  ...
  return $exit_code || true
```

The `|| true` mistakenly swallowed the actual `$exit_code` returned by `recover.sh`, causing `run_recover()` to always return `0` (success), masking any underlying failures in the wake recovery process.

Fix: The `|| true` statement was removed from the return line:

```bash
  return $exit_code
```

Since `watcher.sh` explicitly omits `set -e` (to prevent pipe breakages from killing the daemon), removing `|| true` safely allows the underlying exit code to propagate without risking daemon termination. The watcher daemon was then reloaded (`launchctl kickstart -k`) to pick up the script changes in memory.

15.5.7 Network Readiness Race Condition on `--restart` (Fixed)

Problem: When running `./toggle.sh --restart`, the script tears down the gateway and immediately begins the startup sequence. On systems with Wi-Fi (or any wireless interface), the OS takes a moment to re-associate with the network after the teardown's DNS/routing changes. If the startup sequence ran before the OS had a default route, `get_active_service` would return nothing, routing bypass targets would be wrong, and DNS/WARP setup would silently fail resulting in a partially configured gateway that self-heals only once the `watcher.sh` re-runs.

Root cause: The original code had no gate before the first network-dependent call (`ACTIVE_SERVICE=$(get_active_service)` at the top of the start branch). This was a pure timing race.

First fix (Wi-Fi specific): A polling loop on `ipconfig getifaddr en0` was added to wait for an IPv4 address on `en0` before proceeding. This worked but was brittle it only handled Wi-Fi and would fail for Ethernet, USB-C adapters, or iPhone USB tethering.

Generalized fix: Replaced the `en0` check with `route get default | grep 'gateway:'`. This detects a valid default gateway on *any* interface, covering all connection types. The loop polls every 2 seconds with a 60-second hard timeout (then proceeds with a warning rather than hanging forever). The output prints the detected gateway IP so failures are easy to diagnose in logs.

```bash
# In toggle.sh, before ACTIVE_SERVICE=$(get_active_service)
while true; do
  if route get default 2>/dev/null | grep -q 'gateway:'; then
    _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
    echo "  Network ready (default gateway: $_gw)."
    break
  fi
  ...
```

```
done
```

Why not ping? Pinging an IP (e.g. `1.1.1.1`) during restart is unreliable because DNS may still be hijacked from the previous session, and the exit node routing may not yet be cleared, causing the ping to loop back through the tunnel. `route get default` is a pure local kernel query no packets sent.

15.5.8 `caffeinate -i` Doesn't Survive Real Idle Sleep Mesh Devices Blacked Out on Every Cycle (July 18, 2026)

Symptom: User reported mesh devices getting "infuriatingly slow" internet, self-resolving once the Mac stopped sleeping. `pmset -g log` showed **four** real `Idle Sleep` events in a 38-minute span (16:50, 16:54, 16:58, 17:16) while the gateway believed itself protected. Each one suspended [[Colima VM|Colima]] (and with it all 5 containers [[toggle.sh]]'s exit-node/WARP/AdGuard/routing-fix/Tor stack), so every mesh device using this Mac as its Tailscale exit node lost that exit node outright, then had to ride out [[watcher.sh]]'s post-wake recovery (exit-node re-assert, DNS re-hijack, routing flush, container health re-verify) plus the [[Bug - Post-Wake Feedback Loop|45s post-recovery settle cooldown]] before things felt fast again.

Root cause: `start_sleep_prevention()` (`toggle.sh` ~L1133) only ran `caffeinate -i`, which blocks *idle* system sleep but nothing else and per 15.5.1 was already known to not cover *forced* (lid) sleep either. In this session specifically, the gateway had been manually stopped (`toggle.sh` invoked 17:03:59, `phase=stop` 17:05:58) which also kills the `caffeinate` PID via `stop_pidfile_daemon`; before that stop and during the window the gateway *was* running, the four idle-sleep hits still got through, confirming `-i` alone is not a durable guarantee even while "protected."

Fix: Added `-s` alongside `-i` `caffeinate -i -s &` (`toggle.sh` ~L1133). `-s` blocks sleep outright whenever the Mac is on AC power and is a documented no-op on battery, so unplugged runs still sleep normally (no unexpected battery drain) while a charging Mac the common case for a machine acting as a always-on gateway can no longer idle-sleep out from under the mesh. Does not change the lid-close caveat from 15.5.1 (`-s`/`-i` don't override a forced clamshell sleep); that remains best-effort via `watcher.sh` post-wake recovery.

15.5.9 Wi-Fi Auto-Rejoin as an Always-On Listener in `watcher.sh` (July 23, 2026)

Context. With the WARP Watcher no longer false-killing the gateway on a Wi-Fi drop (15.6.2 point 5 / 15.12.25), the remaining half of the user's actual problem was still open: *"sometimes Wi-Fi disconnects this MacBook and if auto-join is off it just sits there without internet."* Pausing prevents the wrong action; it does not perform the right one (bring Wi-Fi back).

Why "remember", not "detect". macOS Darwin 25 redacts the live SSID from every CLI tool (and even the on-disk known-networks store is a data-vault denied to root) unless the process holds Location Services see 15.11.10. The user refused Location outright. So the current network cannot be read by any script. Joining a network **by name**, however, is *not* gated: `networksetup -setairportnetwork "$iface" "$ssid"` associates using the password already in the login Keychain. The design therefore inverts the problem the user remembers their network name(s) once, and we only ever rejoin one of those.

The helper `scripts/wifi-rejoin.sh` (unsigned).

* `set <SSID>` / `forget <SSID>` / `list` manage the remembered list in `.last_wifi_ssid` (one per line, gitignored). Seeded once with the user's primary network name.
* `rejoin` (internal) arms a down-timer (`/tmp/nullexit-wifi-downsince`) and a try-counter on first call; holds a **15s grace** window so a same-AP DHCP blip self-heals before we force anything; then attempts with **progressive backoff** (`302` capped at 300s, tracked via `/tmp/nullexit-wifi-lastjoin`) and **gives up after `MAX_TRIES=6`** per down-episode (drops `/tmp/nullexit-wifi-gaveup`, logs "reconnect manually") so it never hammers a network that can't come up. Loops the remembered list, sends the re-associate request for the first one that's accepted. It also powers the radio on first if `networksetup -getairportpower` reports it off.
* `reset` (internal) clears all four markers once the link is confirmed back.

Where the driver lives and why that mattered. The first implementation drove `rejoin` from the WARP Watcher (`toggle.sh`) only. That watcher exists **solely while the gateway is up**, so a Wi-Fi drop while the gateway was down or off left nothing to reconnect the Mac confirmed live when a retest "did nothing" because the gateway had shut down and taken its watcher with it (15.12.25, part 3). The fix moved the driver into the **always-on** `watcher.sh` LaunchAgent daemon (`com.nullexit.wake-recovery`), which runs independently of gateway state, as a 4th backgrounded listener alongside the wake / network-change / tunnel-fail listeners:

```bash
# Listener 4: Wi-Fi uplink keepalive  polls the Wi-Fi iface every 15s
phys_ip=$(ipconfig getifaddr "$wifi_iface")
case "$phys_ip" in
  ''|169.254.*) bash "$WIFI_HELPER" rejoin ;;   # no routable IP  try a remembered net
  *)            bash "$WIFI_HELPER" reset  ;;    # link up  clear the down-timer
esac
```

It is leak-safe for the same reason the WARP-Watcher pause is: it only acts when the interface has **no** routable address, so no egress path exists. When the gateway *is* up, the WARP Watcher's pause path calls the same helper too; the two coordinate through the shared `/tmp/nullexit-wifi-*` markers (grace + backoff gate a second caller out) rather than double-associating, and `reset` on either path is idempotent. Verified: with Wi-Fi up the poller stays idle (no markers created); the true end-to-end proof is a manual drop test, which is what confirmed Keychain-based re-association actually works. Cross-refs: 15.5.1 (the daemon and its other listeners), 15.6.2 (the pause guard), 15.11.10 (SSID redaction), 15.12.25 (the

originating incident).

15.6 WARP Tunnel Liveness & Host Leaks

15.6.1 Host-Side Traffic Leaking Past the Exit Node (with Remote Clients Routing Correctly)
Symptom. `toggle.sh` reports success. Remote tailnet clients (e.g. an S24 phone) correctly egress through Cloudflare WARP and see a Cloudflare IP on `whatismyip.com`. But on the **HOST Mac running the gateway**, `whatismyip.com` reports the underlying physical ISP's ASN (e.g. `campus_isp` for a university campus Wi-Fi). The gateway container itself is healthy; the leak is specifically on the host.

Root causes. Three distinct mechanisms can each produce this exact symptom on the host alone, while leaving remote clients unaffected. They are mutually rankable so the diagnostic picks the active one deterministically.

* **(A) SOCKS5 fallback lane active.** `toggle.sh`'s three Phase-B pre-flight checks (`tailscale ping` AdGuard via `dig +tcp` WARP internet, toggle.sh ~lines 750-820) sometimes flake during the first minute. When ANY of the three fails, the script sets `SKIP_EXIT_NODE=true` (toggle.sh:849) and instead activates the SOCKS5 fallback path (toggle.sh:894). The SOCKS5 fallback hijacks DNS to `127.0.0.1` (so AdGuard filtering still works) but **modern browsers on macOS do NOT honor the system SOCKS5 proxy** they ignore `networksetup -setsocksfirewallproxy`. Result: DNS gets ad-filtered but actual HTTP/S traffic exits direct through the campus ISP. This branch is invisible from a remote tailnet client because the client's tunnel is unaffected.

* **(B) IPv6 leak over campus Wi-Fi.** The Tailscale exit node and WARP tunnel are both IPv4-only (gluetun + `post-rules.txt` drop all IPv6 forwarded traffic; see 15.12 IPv6 Exit-Node Leak). But many university and enterprise APs are dual-stack the host happily accepts AAAA DNS responses and sends IPv6 traffic straight out `en0`. macOS happily encrypts that traffic through Tailscale ONLY if Tailscale has IPv6 advertised; with an IPv4-only exit node, IPv6 traffic skips Tailscale entirely. This branch is invisible from a phone because iOS/Android Tailscale apps actively sinkhole IPv6 in their VPN profile, but the macOS host's `tailscaled` does not.

* **(C) Route-table freeze and `--accept-routes` reset after `tailscaled` wake/toggle.** Running `tailscale up --reset` clears the exit-node preference and reverts `--accept-routes` back to its default (`false`). Without explicitly passing `--accept-routes=true`, `tailscaled` will not attempt to install the default route through the `utun*` tunnel interface even if the exit node preference is reconciled. Additionally, macOS occasionally freezes and fails to apply route-table changes (see 15.11 Standalone Tailscale Daemon Freeze). In both cases, `netstat -rn` shows the default route still pointing to the physical Wi-Fi gateway (e.g. `172.17.0.1`) instead of the Tailscale `utun*` interface, causing host traffic to exit direct while remote clients route correctly.

Why remote clients don't see this. Each remote phone/laptop uses its OWN Tailscale tunnel to reach the container's `100.x.x.x:443` over the mesh they're end-to-end encrypted regardless of how the host Mac itself is configured. Only the HOST's own egress sockets (HTTP requests from the host's browser, curl, etc.) are affected.

Resolution. A single script: `scripts/diagnose-host-leak.sh`. It runs the 8 checks, classifies into one of A / B / C / OK, prints a verdict + a ready-to-run fix command on stdout, AND writes the full report to `host-leak-diagnostic-<UTC>.txt` so the user can paste the file back instead of scrolling-and-copying from the terminal. With `--fix`, the matched remediation is applied and the two checks that could change (warp + ipv6) are re-verified. With `--watch` (or `--watch 30`), it runs the full baseline diagnostic then enters a continuous loop checking warp/IPv6/default-route every N seconds, alerting only on state changes logs to `host-leak-watch.log`.

```bash
# Diagnose only:
bash scripts/diagnose-host-leak.sh

# Diagnose + apply matched fix + re-verify:
bash scripts/diagnose-host-leak.sh --fix
```

What it prints. A 7-section report (`tailscale status`, SOCKS5 state, `cdn-cgi/trace` warp=, IPv6 leak probe, host default route, last toggle.sh output.log hints, resolved gateway Tailscale IP), followed by a classification block (`Scenario: A | B | C | OK`) with the exact command(s) the user needs to paste into their terminal they don't have to figure out which `tailscale up` flags match their environment.

What it does NOT do. It does not touch `toggle.sh` or `toggle.sh`'s pre-flight logic, does not modify `.env` (except `--fix` mode appends `GATEWAY_BYPASS_PING=true` when fixing Scenario A), and does not restart Docker. It is a read-only diagnostic with an opt-in remediation flag. Safe to run on a healthy gateway the worst case is a 30-line report that says "OK".

Why cdn-cgi/trace over whatismyip.com for the verdict. `whatismyip.com`'s ISP database routinely mislabels campus IPv6 ranges as residential ISPs (that's why the local ISP's ASN appears even when you're routing through Cloudflare). `cdn-cgi/trace` is hosted by Cloudflare ITSELF and reports `warp=on|off` plus the exact edge colo serving the request. It is the only public endpoint that can truthfully confirm whether the WARP tunnel is in use.

Common false alarms.

* **`warp=on` but whatismyip.com still shows the local ISP.** `whatismyip.com`'s ISP-detection database is unreliable on academic networks. Trust `cdn-cgi/trace`'s `warp=on` over `whatismyip.com`'s ISP string. Use

`https://api.ipify.org` (returns just the IP, no ISP DB lookup) as a third confirmatory check.
* **`tailscale status` shows the host online but exit node preference is missing.** `tailscale up --reset --exit-node=` (empty value) clears any stale preference; `tailscale up --reset --exit-node=<100.x.x.x>` re-establishes it. The diagnostic prints the exact command with the resolved TS_IP filled in.
* **Both IPv6 leak AND route-freeze flags set simultaneously.** Scenario B outranks Scenario C fixing IPv6 makes IPv4 the only viable egress, after which the route-freeze usually self-resolves within ~30s (or one more `toggle.sh` cycle).

Deep dive: Host Leak Probe Continuous Sub-Second Egress Monitor (`scripts/host-leak-probe.sh`).

> ** STATUS: REMOVED (historical).** `scripts/host-leak-probe.sh` and the `HOST_LEAK_PROBE` env flag no longer exist the file is deleted/untracked and nothing launches it (verified 2026-07-15). Its fail-closed role was superseded by the **PF kill-switch** (`enable_killswitch` in `common.sh`) plus the in-container **WARP Watcher** (15.6.2); the only surviving host-leak tool is the on-demand `scripts/diagnose-host-leak.sh`. The text below is kept for design rationale only treat every present-tense claim (auto-launch, PID file, `HOST_LEAK_PROBE=false` toggle) as past tense.

What it is. `scripts/host-leak-probe.sh` is a lightweight background daemon (~1.4 MB RSS, ~01% CPU) that polls `https://www.cloudflare.com/cdn-cgi/trace` every 300ms directly from the host via `curl` not via `docker compose exec`. This is a fundamentally different vantage point from the WARP Watcher (15.6.2): the WARP Watcher asks the WARP *container* whether it sees `warp=on`; the Host Leak Probe asks what a real browser request leaving the host's physical NIC looks like. The two blind spots it covers that the WARP Watcher cannot:

1. **Host-side flash-leaks** a Cloudflare anycast edge re-route, a browser reusing a stale pooled connection across a tunnel state change, or a split-second routing gap during a Gluetun healthcheck restart (see 15.12.13 Routing Fix 30-second polling acknowledged 14s window).
2. **Host egress while container is healthy** the container can report `warp=on` while the host's own traffic exits direct (the three-scenario leak class documented above in 15.6.1).

Lifecycle. Started by `toggle.sh`'s START path (`start_host_leak_probe`) immediately after `start_warp_watcher`. Controlled by:

Signal path	What happens
`toggle.sh` STOP	`stop_host_leak_probe()` SIGTERM + PID file cleanup
`toggle.sh` `cleanup_handler` (error/INT/TERM/HUP)	Same
`recover.sh` (default, non `--post-wake`)	6d block kill by PID file
`recover.sh --post-wake`	Left running the gateway is still live

PID file: `/tmp/nullexit-host-leak-probe.pid`. Toggle with `HOST_LEAK_PROBE=false` in `.env` to disable at next gateway start.

Output format. All events are written to `output.log` (repo root) with UTC timestamps. Only state *changes* are logged the probe is silent during normal healthy operation.

```
[09:10:26] LEAK warp=off ip=45.23.1.1 prev_warp=on prev_ip=104.28.246.50
[09:10:26] ROTATE warp=on ip=104.28.248.11 prev_ip=104.28.246.50
[09:10:26] HOST-PROBE failed/timeout (count=1)
```

curl transport errors (`TLS handshake failed`, `Connection refused`, etc.) are also appended to `output.log` via `2>>output.log` nothing is silenced to `/dev/null`.

Grep cheatsheet.

```bash
grep LEAK output.log            # real warp=off events from the host
grep ROTATE output.log          # Cloudflare anycast edge rotations (harmless)
grep 'HOST-PROBE' output.log    # curl failures / timeouts
```

Resource budget. ~1.4 MB RSS, ~0% CPU at rest, ~3 HTTPS requests/second to Cloudflare. Safe to leave running for hours. Back off the polling interval (`bash scripts/host-leak-probe.sh 1.0`) if you observe probe-failure pile-ups indicating Cloudflare rate-limiting.

Detection vs prevention. This script detects leaks; it does not prevent them. A sub-300ms flash-leak can still slip between two polls undetected. If `grep LEAK output.log` shows confirmed leak events, the next step is a kill-switch (prevention) an `iptables` rule on the host's `en0` that drops non-Tailscale, non-local egress whenever the gateway is active. That is a separate engineering task not yet implemented.

15.6.2 In-Flight WARP Tunnel Liveness Monitor & Statistical Auto-Shutdown
Why. nullexit's core promise is that your IP always appears as Cloudflare. If the WARP tunnel silently dies due to a UDP NAT timeout, a Cloudflare edge rotation, a Colima VM clock skew causing TLS cert rejection, or any of a dozen other failure modes the host silently falls back to egressing through the physical ISP. The user sees the gateway as "up" (containers running, Tailscale connected, DNS hijacked) but their real IP is exposed. This is the exact class of failure that Ingo Blechschmidt documented in his chilling 2024 Linux kernel post-mortem: for over two years, LUKS disk encryption on suspend was silently broken because a refactored kernel syscall returned success while doing nothing ([source](https://mathstodon.xyz/@iblech/116769502749142438)). The lesson: **a security mechanism that fails silently is worse than no mechanism at

```
      all**  it creates a false sense of safety that prevents the user from taking corrective action.

**Design principle.** The WARP Watcher (`start_warp_watcher()` in `toggle.sh`) is a background daemon that polls `cdn-cgi/trace` every 30 seconds and applies a **statistically calibrated threshold** before triggering any safety response. This threshold distinguishes transient network blips (which self-heal within seconds) from genuine outages (which require intervention). A single `warp=off` reading is meaningless  Docker might be slow, the network might be congested, a container healthcheck might be restarting. But 6 consecutive off readings (30 continuous seconds of downtime) has vanishingly low probability of being a false positive: even at an unrealistically high 10% false-positive rate per poll, P (6 consecutive) = 0.1  0.0001%. In practice, the per-poll false-positive rate is far lower than 10%, making the real probability astronomically small.

**What it does.**

1. **Always logs.** Every state transition (onoff, offon) is timestamped to `output.log`. The user can `grep WARP output.log` to audit the tunnel's entire lifetime. This is the "never fail silently" mandate.

2. **Tracks consecutive failures.** A counter increments on each `off` reading and resets to 0 on any `on` reading. This ensures the watcher never fires on intermittent flapping.

3. **Auto-shuts down at threshold.** When the counter reaches `WARP_FAIL_THRESHOLD` (default 6, configurable in `.env`), the watcher declares a statistically significant outage and runs `recover.sh`  the nuclear recovery script that disconnects Tailscale, stops Docker containers, resets DNS to `1.1.1.1`, flushes routes, and power-cycles Wi-Fi. This instantly reverts the host to normal (un-encrypted) internet, guaranteeing no traffic leaks through a dead tunnel while the user thinks they're protected. A trigger file at `/tmp/nullexit-warp-shutdown-triggered` prevents double-firing.

4. **Exits cleanly.** After shutdown, the watcher exits (the gateway is dead, there's nothing left to monitor). `stop_warp_watcher()` is also called by `toggle.sh`'s STOP path and `cleanup_handler` to terminate the daemon cleanly during normal teardown.

5. **Pauses instead of killing when the *physical uplink* is down (added 2026-07-23).** The watcher cannot tell "WARP tunnel failed" apart from "the Mac lost Wi-Fi"  both read `warp=off`, because with no internet there is no tunnel either. Before this guard, a plain Wi-Fi drop counted 6 `off` polls and nuked the whole gateway (the 2026-07-22 outage, 15.5.9). Now, before counting a `warp=off` as a failure, the loop resolves the physical interface (`networksetup -listallhardwareports`, same as `wait_for_dhcp_settle`) and checks `ipconfig getifaddr $phys_iface`. If it's empty or `169.254.*` (self-assigned) there is **no routable uplink**, so it logs `WARP watcher PAUSED`, resets `consec_off=0`, calls `scripts/wifi-rejoin.sh rejoin`, sleeps, and `continue`s  it does **not** run `recover.sh`. When the interface gets an IP again it logs `RESUMED` and calls `wifi-rejoin.sh reset`. This is **leak-safe**: the pause only ever fires when the interface has no routable address, i.e. there is no egress path at all and therefore nothing that *can* leak. The kill-switch semantics are unchanged for the case that actually matters  "link up **and** `warp=off`" still trips the nuclear shutdown. Note the pause covers the gateway-up case; the always-on `watcher.sh` (15.5.9) is what drives rejoin when the gateway is down.

   **Link-state refinement (added 2026-07-23, after the guard's first version let the false-kill recur 15.12.27).** The original `getifaddr`-only check has a hole: *having an IP is not the same as having egress.* On an enterprise-Wi-Fi (802.1x) de-auth or a hard roam, macOS keeps the **stale DHCP lease** on `en0` for a while after the association is already dead, so `ipconfig getifaddr en0` still returns a real, non-`169.254` address. The bare check saw "uplink present," skipped the pause, the container's `cdn-cgi/trace` probe legitimately failed (no real internet), 6 `off` polls accumulated in ~50 s, and the gateway got nuked anyway  the exact failure this guard was supposed to stop. So after confirming a routable IP, the watcher now also reads the **driver's link-layer association state**  `ifconfig $phys_iface | awk '/status:/{print $2}'`. `status: active` = associated (proceed to the WARP check); `inactive` = link down under a stale lease  treat as no uplink and PAUSE. **Why not ping the router?** The obvious "does the next hop answer?" probe is *wrong here*: nullexit's own PF kill-switch plus this network's AP client-isolation mean the LAN gateway drops ICMP **even when the tunnel is perfectly healthy**  verified live, a ping to `172.17.0.1` fails while `warp=on`. A ping-based guard would therefore false-PAUSE a live gateway, which is worse than the bug it fixes. The link-state read is ICMP- and PF-independent. It **fails open**: any status other than a definitive `inactive`/`none` (including an absent status line) trusts the lease rather than pausing, so an unexpected driver string can't strand the watcher. Fail-open is safe because the PF kill-switch remains the actual leak guarantee  a wrongly-paused watcher yields "no internet," never "real IP exposed." The `PAUSED`/`RESUMED` log lines now read "no working uplink" / "working uplink back."

**Configuration.** Set `WARP_FAIL_THRESHOLD=3` in `.env` for a more aggressive 15s timeout, or `WARP_FAIL_THRESHOLD=12` for a conservative 60s window. The variable is parsed from `.env` using the same `grep` pattern as `GATEWAY_MSS` and `GATEWAY_BYPASS_PING`.

**Log sample.** After a real outage (or a forced test via `docker compose pause warp`):

```
[2026-07-03T21:15:17Z] WARP DOWN  cdn-cgi/trace reports warp=off
[2026-07-03T21:15:47Z] WARP SHUTDOWN  6 consecutive failures (threshold=6). Running recover.sh to kill the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.
[2026-07-03T21:15:52Z] WARP SHUTDOWN  recover.sh completed, watcher exiting. Your IP is now your real ISP IP.
```

Or, if WARP self-recovers before the threshold:

```
[2026-07-03T21:15:17Z] WARP DOWN  cdn-cgi/trace reports warp=off
```
```

```
[2026-07-03T21:15:32Z] WARP RECOVERED  warp=on after 3 consecutive off readings (threshold was 6)
```

Relationship to other monitors. The WARP Watcher is the **innermost safety layer** it runs as a child of `toggle.sh` and only while the gateway is active. It complements (does not replace) the `scripts/watcher.sh` daemon (15.5.1, which handles sleep/wake and Wi-Fi roam) and `scripts/diagnose-host-leak.sh` (15.6.1, which is a manual diagnostic tool). The WARP Watcher is the only component that actively verifies end-to-end tunnel liveness and takes autonomous action when it fails.

15.6.3 Incident: The Overnight Silent IP Leak (July 2026)
The Problem: The host leaked traffic over its raw, unencrypted `eth0` IP continuously for roughly 8 hours, completely bypassing the VPN. The `toggle.sh` state and the container liveness checks all reported the gateway as healthy and active.

The Investigation:
1. **Sleep/Wake Checks:** We initially suspected macOS sleep cycles broke the routing table. A check of `pmset -g log` confirmed the laptop literally never slept (0 sleep events recorded since boot), ruling out all sleep/wake code paths.
2. **Keepalive Expiry:** We discovered `docker-compose.yml` was missing `WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL`. Without a keepalive, standard router NAT tables (which typically garbage collect UDP after 30 to 120 seconds of silence) quietly expired the WireGuard port mapping overnight.
3. **Container State Masking:** Because WireGuard is stateless, the `warp` container didn't immediately crash. It stayed "Up", which meant Docker's healthchecks and our `toggle.sh` liveness checks never noticed the tunnel inside it was totally dead.

The Solution:
1. Added `WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL=25s` (the WireGuard standard to beat standard 30s NAT timeouts) to keep the UDP tunnel permanently pinned open through the router.
2. Injected a **Fail-Closed Kill-Switch** into `routing-fix.sh`: A 30-second loop now actively verifies tunnel health by running `curl` over `tun0` to Cloudflare's trace endpoint. If it fails 3 consecutive times, it immediately rewrites the default route to `blackhole`, severing all traffic instead of letting it leak through `eth0`.
3. Added a listener in `watcher.sh` that detects this fail-closed event and instantly pushes a native macOS UI notification to the desktop, alerting the user to manually intervene.

> [!NOTE]
> **Why dual failure systems?**
> The system now relies on two independent failure handlers that cover each other's blind spots:
> 1. **Automated Self-Healing (WARP Watcher + `recover.sh`):** Triggered when the `warp` container logs `warp=off`. This happens when the server actively kicks the client. Because the container *knows* it is disconnected, it triggers `recover.sh` to safely reboot Docker and reconnect.
> 2. **Fail-Closed Kill-Switch (`routing-fix.sh` + `watcher.sh`):** Triggered when a silent network failure occurs (e.g., NAT timeouts). The container incorrectly believes it is connected (`warp=on`), so auto-recovery never fires. Instead of auto-recovering (which risks falling back to raw Wi-Fi mid-process while the user isn't looking), this kill-switch actively blackholes the internet and forces a manual intervention via a desktop notification.

15.7 Tailscale P2P, SNAT & Control Plane

15.7.1 macOS SNAT Endpoint Poisoning & The P2P Blackhole (July 10, 2026)
Symptom: When a Tailscale client (like an S24) connected to the exact same Wi-Fi network as the macOS gateway, it completely lost internet connectivity. However, if the S24 connected to a Windows PC's hotspot on the same network, the tunnel worked flawlessly.

Root Cause: Claude's critique correctly dismantled the port collision theory: the gateway daemon wasn't failing to bind, and macOS wasn't load-balancing ports. The real culprit was **macOS SNAT Source Port Mangling poisoning Tailscale's path discovery**.

Here is the exact step-by-step of the "infinite loop" blackhole:
1. **The P2P Handshake:** The S24 and the Mac are on the same Wi-Fi (`192.168.1.x`). The S24 discovers the Mac's local IP and sends a UDP hole-punch packet to `192.168.1.5:41641`. Colima successfully forwards this to the gateway container.
2. **The Bypass Rule:** The gateway replies. Because `routing-fix.sh` forces all Tailscale UDP traffic out `eth0` to bypass WARP, the reply goes to the macOS host.
3. **macOS SNAT Mangling:** macOS routes the reply out `en0` to the S24. Because the packet crosses from the Colima VM bridge to the Wi-Fi interface, macOS applies Source NAT (SNAT). Critically, macOS **randomizes the source port** (e.g., changes it from `41641` to `58291`).
4. **Endpoint Poisoning:** The S24 receives the authenticated Tailscale packet from `192.168.1.5:58291`. Tailscale's `magicsock` is designed to handle NAT dynamically, so it assumes the Mac has roamed to port `58291`! The S24 updates its endpoint and starts sending all its encrypted internet traffic to `192.168.1.5:58291`.
5. **The Blackhole:** macOS has no port forward for `58291`. It drops 100% of the S24's outbound data packets. The connection is completely dead.

Why did the Hotspot work? When the S24 connected to the PC hotspot, it was behind the PC's NAT. The S24 sent the packet, PC NATted it, and the Mac replied (mangled to `58291`). When the reply hit the PC, the PC's strict NAT dropped it because the source port (`58291`) didn't match the destination port it originally sent to (`41641`). Because the mangled packet was dropped, the S24 never poisoned its endpoint! It safely gave up on P2P, fell back to the 100% reliable TCP DERP relay, and the internet worked perfectly.

Fix & Resolution (July 10, 2026):

Phase 1 — Port disambiguation: Changed the gateway container's Tailscale listen port from the default `41641` to `41642` across `docker-compose.yml`, `post-rules.txt`, and `routing-fix.sh`. This prevents any bind conflict with the macOS host's native Tailscale daemon.

Phase 2 — RFC1918 DROP rules: Updated `routing-fix.sh`'s `add_tailscale_p2p_bypass()` to explicitly drop all outgoing Tailscale UDP packets destined for private subnets (`192.168.0.0/16`, `10.0.0.0/8`, `172.16.0.0/12`) when `TAILSCALE_ALLOW_LAN_P2P=false`. This surgically kills the local P2P hole-punch attempt *before* it can trigger the macOS gvproxy SNAT mangling. The S24 receives a clean failure and immediately locks onto the public DERP relay with 0% packet loss.

Phase 3 — Automatic network detection: Because `TAILSCALE_ALLOW_LAN_P2P=false` breaks P2P on trusted hotspots that genuinely don't have AP isolation, a `.env` toggle and a full auto-detection system were implemented:

- **`.env` toggle:** `TAILSCALE_ALLOW_LAN_P2P=true/false` — static fallback, used only if auto-detection cannot run.
- **`watcher.sh` auto-detection (`detect_lan_p2p_mode`):** On every network change, `watcher.sh` runs two checks:
  1. **WPA2-Enterprise detection:** `airport -I` is used to read the current Wi-Fi `link auth` field. If the auth type does not contain `-psk` (e.g., it is plain `wpa2` = 802.1x), the network is classified as enterprise and P2P is forced off. *(Note: the `airport` binary was removed in macOS Sonoma 14.4 — the detection now uses `system_profiler SPAirPortDataType`; see 15.11 for the full migration.)*
  2. **AP Isolation probe:** For WPA2-Personal networks, a short `ping` is sent to the default gateway. If the gateway responds (it always should on a real home/hotspot network), P2P is allowed. If not, AP isolation is inferred and P2P is forced off.
  - The result (`true` or `false`) is written to `.lan_p2p_detected` in the repo root.
- **`routing-fix.sh` dynamic reading:** `add_tailscale_p2p_bypass()` re-reads `.lan_p2p_detected` on every 30-second loop iteration and atomically adds or removes the RFC1918 DROP rules. **No container restart is needed when switching networks.**

Known Unknowns (Pending Verification):
- The gvproxy SNAT source port mangling theory is mechanistically sound and backed by observed behavior (gvproxy's UDP proxy changelog has documented edge cases in this exact code path), but has not yet been confirmed via `tcpdump -i en0 udp portrange 40000-60000` while triggering a P2P handshake from the phone. A packet capture cross-referenced with `tailscale ping` endpoint reports on the phone side would confirm the theory definitively.
- Claude's recommendation: run `tcpdump -i en0 udp port 41642 or portrange 40000-60000` on the Mac while triggering the handshake from the phone to see whether the reply leaves with a different source port than `41642`.
- **Client-Side VPN Cycling (Roaming Recovery):** An observation (July 10, 2026) suggests that when a hung DNS probe state is triggered by roaming between Access Points (APs) on a WPA2-Enterprise network, simply toggling the VPN connection (Tailscale) off and on locally on the client phone immediately resolves the issue. You do not need to cycle the phone's physical Wi-Fi. This reinforces the theory that roaming on Enterprise networks leaves a stale/poisoned P2P endpoint in the phone's Tailscale magicsock state, which gets instantly flushed upon a local VPN restart. (Status: Possible but not formally verified).

15.7.2 July 6, 2026: Packet Filter (pfctl) Defaulting Local LAN Rules to TCP-Only
Symptom: Devices on the exact same local Wi-Fi network as the gateway were experiencing ~500ms latency to the gateway instead of the expected ~80ms.

Root Cause: When writing the `pf.conf` rules to whitelist local LAN traffic (`pass out quick on $ext_if to { 192.168.0.0/16, ... }`), the rule omitted an explicit protocol. The macOS `pfctl` compiler implicitly applies state tracking to all `pass` rules. However, because state tracking (`keep state`) natively evaluates TCP sessions, `pfctl` automatically appended the `flags S/SA` modifier to the compiled rule in the kernel. This silently restricted the whitelist to **TCP only**, causing all local UDP traffic to be dropped by the default-deny kill-switch. Since Tailscale relies on UDP port 41641 for direct P2P connections, the local devices were unable to hole-punch and were forced to fall back to routing traffic through a remote Tailscale DERP relay, massively inflating latency.

Resolution: Split the single implicit rule into explicit `proto tcp`, `proto udp`, and `proto icmp` rules in `scripts/pf.conf` to guarantee the macOS compiler allows all three core transport protocols on the local subnet without mutating the rule into a TCP-only lock.

15.7.3 Incident Post-Mortem: PF Kill-Switch Bricks Host Internet in SOCKS5 Fallback Mode (July 10, 2026)
Symptom: When the pre-flight checks fail (e.g. because host Tailscale was unauthenticated/logged out), the script falls back to SOCKS5 proxy mode but the host immediately loses all internet connectivity and tailscaled reports `You are logged out... connect: no route to host`. Even running `sudo tailscale up` to log in fails because the connection to the control plane times out.

Root Cause: An architectural mismatch. The macOS Packet Filter (PF) Kill-Switch works at the IP level: it blocks all outbound traffic on the physical interface (`en0`) except to explicitly whitelisted VPN endpoints, expecting all other traffic to go through the virtual VPN interface (`utun*`). In SOCKS5 fallback mode, the exit-node (`utun*`) is NOT enabled; instead, only specific applications (like browsers) use the localhost SOCKS5 proxy, while all other host traffic (including the `tailscaled` daemon's UDP packets) continues to go through `en0`. Because the script still enabled the PF Kill-Switch during SOCKS5 fallback, it blocked all non-SOCKS5 traffic on `en0`, completely bricking the host's internet and locking the Tailscale daemon out of the control plane.

Fix (July 10, 2026): Modified `toggle.sh` to remove the `enable_killswitch` call from the SOCKS5 fallback path. The firewall kill-switch is now strictly reserved for exit-node VPN mode, ensuring SOCKS5 fallback leaves the host's direct internet route open for system daemons to authenticate.

15.7.4 Incident Post-Mortem: Host Tailscale Logouts due to Aggressive reset Flags (July 10, 2026)
Symptom: Every time the user runs `toggle.sh` or `recover.sh` which disconnects/reconnects the host client, the host client randomly gets logged out, forcing them to re-run `sudo tailscale up` to authenticate.

Root Cause: An overly aggressive configuration reset. The host-side `tailscale up` calls (used in `disconnect_tailscale_host` and the startup/enable exit node paths) all passed the `--reset` flag. On macOS, `--reset` resets the entire configuration state and authentication token cache of the daemon. Since the script does not pass a `TS_AUTHKEY` to the host's commands (which are meant for the user's private interactive client), the `--reset` flag effectively logged the user out of their tailnet, requiring interactive re-login.

Fix (July 10, 2026): Removed the `--reset` flag from all host-side `tailscale up` invocations in `toggle.sh` and `scripts/common.sh`. This ensures that exit-node state transitions (enabling/disabling) are handled safely while preserving the host client's authenticated session.

15.7.5 Incident Post-Mortem: macOS Tailscale Flag Requirement Abort (July 10, 2026)
Symptom: When the `toggle.sh` stop trap or the roaming watcher triggered, the gateway failed to shut down or disconnect properly. The `output.log` showed the error: `Error: changing settings via 'tailscale up' requires mentioning all non-default flags. To proceed, either re-run your command with --reset...`

Root Cause: In a previous attempt to stop host-side logouts (15.7.4), we removed the `--reset` flag from all `tailscale up` invocations. However, if the Tailscale daemon on macOS is already running with non-default flags (like `--ssh` or `--accept-routes`), invoking `tailscale up` with only a subset of flags (like `--exit-node=`) triggers a strict safety check in the Tailscale CLI that aborts the command completely. Because the command aborted, exit-node routing remained stuck.

Fix (July 10, 2026): Migrated all specific preference changes in `toggle.sh` and `scripts/common.sh` from `tailscale up` to a two-step sequence:
1. `tailscale up` (with no flags) to safely ensure the interface is brought up using the current authenticated state.
2. `tailscale set --exit-node=...` to modify the specific target preference cleanly without triggering the missing flags error or requiring `--reset`.

15.8 Docker / Colima Lifecycle

15.8.1 Colima VM Memory / Swap Thrashing Latency
Symptom: After running nullexit flawlessly for over 24 hours straight to test stability, the 0.5GB VM RAM configuration began experiencing severe network latency on the Tailscale mesh later in the day.

Root Cause: Services (like AdGuard, Tailscale, Docker logs) slowly accumulate cache and state over extended periods. Under the tight 512MB physical RAM limit, this slow creep forced the Linux kernel to aggressively swap memory to the 512MB SSD swap file. The constant disk I/O thrashing blocked process execution, causing DNS resolutions and packet forwarding to delay significantly (making the internet feel "very slow").

Proof (Empirical Observation): While in this OOM-thrashing state, direct DNS queries through the Tailscale mesh (`dig @100.100.21.8 google.com`) would intermittently hang or take several seconds to resolve. After restarting with the 600MB configuration, the exact same query immediately dropped to a lightning-fast **41 ms** response time. (Note: Querying the mapped host port via `127.0.0.1:5354` will always time out due to Gluetun's strict leak-prevention firewall blocking non-VPN ingress traffic).

Fix: Increased VM base physical RAM to 600MB (`--memory 0.6`) and reduced the swap file to 400MB. This provides enough native RAM headroom for long-running state caches without forcing heavy I/O swapping. Subdomain deduplication still reduces blocklists by ~60%. Memory profiles (`light`/`medium`/`heavy`) let users trade off coverage for memory.

15.8.2 Missing iptables in routing-fix Container
Root Cause: `alpine:3.20` does not have iptables installed. Commands failed silently with `2>/dev/null || true`.

Fix: Updated `docker-compose.yml` to `apk add --no-cache iptables iproute2` before `scripts/routing-fix.sh`.

15.8.3 DOCKER_GW Variable Parsing Bug
Root Cause: `ip route show default` printed two lines; `awk '{print $3}'` picked up both, causing route commands to fail.

Fix: Added `head -1` to the parsing chain.

15.8.4 Hard Reboots Leave Stale Docker Socket in Colima VM
Problem: If the macOS host machine is subjected to a hard reboot, sudden power loss, or a kernel panic, the Colima VM running in the background is abruptly killed. Upon next boot, running `toggle.sh` resulted in the contradictory output:
`Colima is already running.` followed by `Cannot connect to the Docker daemon at unix://~/.colima/default/docker.sock.`

Root Cause: The hard crash prevents the inner Docker daemon from executing its graceful shutdown routines. It leaves behind a corrupted `docker.sock` file inside the VM. On boot, the outer VM spins up (hence "

```
     already running"), but the inner Docker daemon sees the stale socket, assumes another instance is running,
     and crashes.
1550
1551 **Fix:** Added an Auto-Healing routine directly into `toggle.sh` (Step 4b). The script now explicitly tests the
      Docker daemon with `docker ps`. If the daemon is unresponsive despite Colima reporting as active, it
     assumes a stale socket and automatically runs `colima restart` in the background to cleanly rebuild the VM
     and socket before proceeding.
1552
1553 #### 15.8.5 Cloudflare WARP 24-Second Startup Delay (UDP Session Rate-Limiting)
1554 **Context:** When running `toggle.sh` to turn the gateway ON, `docker compose up -d` often hangs for exactly
     24-25 seconds waiting for the `warp` container to become `healthy`. Users might mistakenly blame the Go `
     rule-compiler` processing 400k+ rules for this delay, as Docker prints `rule-compiler Exited 24.3s` at the
     end of the wait.
1555 **Root Cause:** The `rule-compiler` actually finishes its work in <2 seconds. The 24-second blockage is
     entirely due to Cloudflare WARP's edge server rate-limiting. WireGuard is a stateless protocol with no "
     disconnect" packet. When the toggle is turned OFF, the old UDP session state remains active on Cloudflare's
      backend for 20-30 seconds. When the toggle is instantly turned back ON, Docker starts a new WireGuard
     connection from a new ephemeral UDP source port but uses the **same static cryptographic key** (from `.env
     `). Cloudflare detects this as a potential replay attack or key-sharing abuse, and intentionally drops all
     packets (rate-limits) for the new connection until the old session's ghost state fully expires (~20-25
     seconds).
1556 **Result:** Gluetun's healthcheck fails repeatedly every 2 seconds until Cloudflare lifts the penalty box, at
     which point traffic flows and Docker unblocks the rest of the startup sequence. This is a known, expected
     behavior of reusing static keys rapidly on Cloudflare's network and cannot be bypassed.
1557
1558 #### 15.8.6 July 4, 2026: Multi-stage Docker Builds for Go Ports & Teardown Hang Fix
1559 **Symptom:** The Python implementations of `logger.py` and `sync-rules.py` were slow and required a heavy
     python runtime inside Alpine containers. Also, `routing-fix` container teardown was hanging for 30s during
     `toggle.sh` shutdown.
1560 **Root Cause:**
1561 - Python is slower than Go and installing it dynamically via `apk add` added overhead and complexity to the
     containers.
1562 - Docker containers ignore `SIGTERM` if the init process doesn't handle it, causing a full 30s timeout on
     shutdown.
1563 **Resolution:**
1564 - Ported `logger` and `sync-rules` to Go using Goroutines for massive concurrent speedups.
1565 - Implemented **Multi-stage Docker builds** (using `golang:1.22-alpine` as builder) to compile the static
     binaries inside Docker. The final containers are raw Alpine with zero dependencies.
1566 - Added `--build` to `docker compose up -d` in `toggle.sh` to ensure updates are consistently applied on boot.
1567 - Added `stop_grace_period: 1s` to `routing-fix` in `docker-compose.yml` which eliminates the 30-second
     teardown hang instantly.
1568 - Implemented a cryptographic integrity checker (`scripts/crypto.sh`) using HMAC-SHA256 and `NULLEXIT_SEED` in
      `.env` to prevent tampering of bash scripts. (See also 15.9.4.)
1569
1570 #### 15.8.7 Colima `start` Post-Provision Hang (~2 min) & the Docker-Readiness Poll (July 14, 2026)
1571 **Symptom:** `toggle.sh` took ~197s of wall-clock time to bring the gateway up, and almost the entire delay
     lived in a single step: `colima start`. Timestamps in `output.log` showed the VM reaching `READY` and
     creating the `colima` Docker context in ~8-12 seconds, after which `colima start` simply *did not return*
     for ~110 more seconds  until the `run_with_timeout 120` watchdog `SIGKILL`ed it. The rest of the boot (
     containers, routing) then proceeded normally.
1572
1573 **Root Cause:** A colima-internal hang that occurs *after* `Successfully created context colima` is printed. It
      is **mode-independent**. It was initially misattributed to `--network-address` / `--network-mode bridged`
     needing the `socket_vmnet` helper, but a controlled test with plain user-mode networking (`colima start --
     memory 0.6 --vm-type vz`, no `--network-address`) reproduced the identical ~120s hang (04:40:34  04:42:34,
     watchdog-killed) while Docker was fully usable the entire time. Crucially, killing the hung `colima start`
     process does **not** stop the VM and does **not** corrupt state  `colima status` correctly reports `running
     ` afterwards and `docker` keeps working via the `colima` context. The ~110s was therefore pure dead time
     waiting on a CLI that had already done its real work.
1574
1575 **Fix:** `scripts/common.sh` now provides `colima_start_until_ready()`. It launches `colima start "$@"` in the
     background and polls `docker --context colima info` every 3s. The moment Docker answers (~12s in practice)
     it stops waiting, reaps the (usually still-hung) starter with `kill`, runs `docker context use colima` to
     guarantee the default context is set, and returns 0. If Docker never answers within a 90s cap it returns
     non-zero and `toggle.sh` continues to the existing Step 4b `docker ps` auto-heal (15.8.4). Both Colima
     start call sites in `toggle.sh`  the LAN-P2P bridged/shared path and the user-mode fast path  now go
     through this helper instead of `run_with_timeout 120 colima start`.
1576
1577 **Result:** The Colima phase dropped from ~120s to ~12s and total `toggle.sh --restart` wall time from ~197s to
      ~106s, with the double-tunnel and direct hostcontainer P2P confirmed intact after the change. The key
     insight: treat "Docker is reachable"  not "the `colima start` CLI returned"  as the definition of *started
     *.
1578
1579 #### 15.8.8 Toggle-On Startup Optimizations: Rule-Compiler Off the Critical Path (July 14, 2026)
1580 **Context:** A genuine cold toggle-on (VM bounced for RAM, WARP down) measured **95s**. Profiled per-phase with
      the now-fixed `DEBUG_TRACE` (15.11.8-A): Colima cold start **19s** (near floor), `docker compose up` **45s
     **, tailscale connect+poll **21s**, final checks **7s**.
1581
1582 **Two wrong hypotheses first (documented so they aren't re-tried):** (1) that BuildKit `--build` was the ~30s
     cost  it wasn't; every layer was already `CACHED`, so gating `--build` saved only ~2s. (2) that `docker
     compose exec` was slow in the poll loop  it isn't (0.18s steady-state vs 0.06s for `docker exec`); the ~5s/
```

iteration was `tailscale status` **blocking during cold connect** and hitting the `run_with_timeout 5` cap.

Real bottleneck (evidence-backed): the 45s `compose up` was **not** WARP. On a genuine cold start WARP goes healthy fast (tor starts right with it no 15.8.5 same-key penalty; that only bites rapid offon). The long pole is the **rule-compiler re-deduping ~340k rules inside the 600MB VM (~28s), every toggle** `adguardhome` started 28s after warp because it waited on `rule-compiler: service_completed_successfully`, even though all 16 remote lists loaded from cache unchanged.

Fix hash-gated, off the critical path (kept in-container):
- `rule-compiler` is now a **`compile`-profiled** service, excluded from `docker compose up`. `adguardhome` and `routing-fix` no longer `depends_on` it they boot instantly on the `compiled_rules.txt` / `ip_blocklist.ipset` already in `./adguard/work`.
- `toggle.sh` runs it on-demand via `docker compose run --build --rm rule-compiler`, **gated on `compute_rules_hash()`** = sha256 of `black_list.txt` + `white_list.txt` (the lists the user controls). Recompile only when that hash changes, when an artifact is missing, or when the compiled file is older than `RULES_MAX_AGE_DAYS` (default 7, so the remote blocklists still refresh occasionally). A compile failure is **non-fatal** we keep the previous compiled rules rather than bricking the tunnel (ad-block the kill-switch).
- Also trimmed the tailscale-connect poll timeouts (status `5s3s`, `ip -4` `10s5s`) to reduce overshoot.

Why NOT run the compiler natively on the Mac host (the tempting idea): rejected on purpose. There is no Go on the host, so "native" means either `brew install go` (a host toolchain dependency that drifts the project from *portable+secure-on-any-device* toward *faster-on-mine*) or committing a prebuilt arch-specific **binary blob** you'd then have to trust/sign. Worse, it would make the **host** write into the `./adguard/work` bind-mount violating the deliberate *"single allowed writer is the rule-compiler container"* invariant (hostcontainer UID/perms mismatches corrupt AdGuard's BoltDB store). Staying in-container preserves both properties; the gate delivers the speed.

Result: unchanged toggle `compose up` **45s 9s**, total `--restart` **~95106s 63s**, with AdGuard still serving the full 342,519 rules (it reuses the existing compiled file) and the double-tunnel + direct hostcontainer P2P intact. Editing `black_list.txt`/`white_list.txt` flips the hash one-off recompile new rules take effect. Both lists were added to the `crypto.sh` signed set (they're security inputs), so editing them requires a re-sign which is the same moment you'd want the recompile. Local state files `.build_hash` / `.rules_hash` are gitignored.

15.9 Permissions, Crypto & Security Hardening

15.9.1 Sudo Credential Caching Removed
Problem: The script used `sudo -v` at startup to cache sudo credentials, plus a background `SUDO_KEEPER_PID` loop refreshing them every 60 seconds. This required interactive password entry even with NOPASSWD sudoers configured, because `sudo -v` itself prompts.

Fix: Removed the entire `sudo -v` + `SUDO_KEEPER_PID` section. All privileged commands now use `sudo -n` (non-interactive), which works silently because the user has NOPASSWD rules in `/etc/sudoers.d/nullexit` for the specific commands needed (`networksetup`, `python3`, `dscacheutil`, `killall`, `pkill`, `route`, `ifconfig`).

15.9.2 Background Sudo Failures
Context: The `--post-wake` script is intentionally designed to skip the `sudo -v` interactive prompt because it is launched by `watcher.sh` running via `launchd` in the background (which has no TTY and cannot accept password input). Thus, it relies entirely on `sudo -n` (non-interactive sudo) for all its internal routing changes.
Problem: If the user has not configured the `/etc/sudoers.d/nullexit` file properly, any automated background wake event will silently fail: `sudo -n` commands drop out, the WARP bypass routes fail to apply, and the `warp` container loses internet connectivity.

15.9.3 Tailscale Hijacking the Default Route (PHYSICAL_GW Bug)
Problem: When `--post-wake` runs, it clears the routing table using `sudo route -n flush` to ensure a clean slate after the Wi-Fi card roams or wakes up. Because it skips bouncing the Wi-Fi interface (to make the reconnect faster), macOS's DHCP client does not automatically rebuild the physical default route. To counteract this, the script must manually restore the physical default route. Originally, the script tried to capture the physical gateway via `route get default`. However, because the gateway is already running and Tailscale is active, the default route is often already pointing to `utunX`. When this happens, `route get default` does not output a `gateway:` field (it only outputs the `interface:`), causing the captured `PHYSICAL_GW` variable to be empty.
When `PHYSICAL_GW` is empty, the physical default route is never restored, causing the `en0` interface to be isolated from the physical network. The WARP bypass routes then fall back to targeting `-interface en0` (instead of routing via the gateway), which breaks WireGuard's remote connections completely.
Fix: The script now bypasses the OS routing table entirely and directly queries the hardware DHCP lease (`ipconfig getpacket en0`) to reliably extract the physical router IP, ensuring the physical route is correctly restored even when Tailscale is active.

15.9.4 July 4, 2026: Consolidation of Core Bash Scripts (Refactoring)
Symptom: Over 200 lines of bash logic were directly duplicated across `toggle.sh` and `recover.sh` (e.g. `restart_tailscaled_daemon`, `stop_sleep_prevention`, WARP bypass routes). This caused drift when one script was updated without the other.
Root Cause: Historical separation of responsibilities allowed `recover.sh` to grow complex alongside `toggle.sh`.
Resolution:
- Massively refactored both scripts by extracting all duplicated networking logic, PID lifecycle management, and environmental configuration down to `scripts/common.sh`.

- Specifically, the `stop_sleep_prevention`, `stop_dns_watcher`, `stop_host_leak_probe`, and `stop_warp_watcher` functions were collapsed into a single `stop_pidfile_daemon()` abstraction.
- Proxy disable loops were abstracted to `disable_all_proxies()`.
- The `162.159.192.1/.193.1` WARP edge IPs are now dynamically resolved from `.env` instead of hardcoded.

15.9.5 Threat Model: Local Proxy Discovery
The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). The real security boundary here is "can an unprivileged local process reach the loopback interface," not "does the malware know the port number."

The port-randomization and the Honey-Port trap only catch blind or generic port-scanners that do not already know where to look. They **do not protect** against a process that directly reads the `/tmp/nullexit-ports.env` file, greps the `toggle.sh` memory space, or runs `lsof -i` to inspect open socket bindings.

Because modifying file ownership on `/tmp/nullexit-ports.env` via `chown`/`chmod` to restrict read-access introduces an unavoidable Time-Of-Check to Time-Of-Use (TOCTOU) race condition (the file is created user-owned before privilege escalation, meaning file descriptors can be snatched), it is deliberately left as a standard user-owned file.

The actual, proper mitigation for preventing a malicious local process from reaching the proxy is not hiding the portit is running a host-level outbound firewall with strict localhost/loopback filtering enabled (e.g., LuLu 4.3.0+ or Little Snitch). These firewalls gate access by cryptographic process identity (binary signature) at the kernel/network-extension level, rather than by port secrecy.

Verifying Localhost Filtering (The Rogue Binary Test). To empirically validate that the host firewall properly intercepts local proxy hijacking, you can compile and execute a completely unknown, un-whitelisted "rogue" binary to attempt a connection to the proxy port:

```c
// lulu-test.c
#include <stdio.h>
#include <stdlib.h>
#include <arpa/inet.h>

int main(int argc, char *argv[]) {
    int port = atoi(argv[1]);
    int sock = socket(AF_INET, SOCK_STREAM, 0);
    struct sockaddr_in server = { .sin_family = AF_INET, .sin_port = htons(port) };
    server.sin_addr.s_addr = inet_addr("127.0.0.1");

    printf("Rogue binary attempting connection to 127.0.0.1:%d...\n", port);
    if (connect(sock, (struct sockaddr *)&server, sizeof(server)) < 0) {
        printf("Connection BLOCKED by firewall.\n");
        return 1;
    }
    printf("Connection SUCCESSFUL (Firewall bypassed!).\n");
    return 0;
}
```
Compile with `gcc -o lulu-test lulu-test.c`. When executed against the `$TOR_SOCKS_PORT`, a properly configured host firewall (with "Allow Loopback" disabled in LuLu Preferences) will immediately pause the socket execution and generate a desktop alert for the unrecognized binary signature. This physically stops targeted local malware from exfiltrating data through the Tor tunnel, proving the threat model mitigation.

15.9.6 Unlocking Mode-Locked Output Files Without `chmod` (`unlock-files.sh`)
The one-command fix: `bash scripts/unlock-files.sh` replaces every known locked inode in the project with a fresh writable one. No `chmod`. Covers `.env` (mode `000`), `adguard/work/userfilters/compiled_rules.txt`, `ip_blocklist.ipset`, `cache/*.txt`, `cache/ip/*.txt`, and `adguard/work/data/filters/*.txt` (mode `0444`). Run it once after setup or after pulling an old repo; `toggle.sh` and `sync-rules.py` will then work without permission errors.

Context: The nullexit project has a strict project-wide policy of `no chmod from scripts` (see README 6). This rule exists because every prior `chmod 0444` (post-write tamper-proof lock) and `chmod 000` (post-toggle `.env` lock) call from `scripts/sync-rules.py` / `toggle.sh` could leave a file permanently unreadable on disk if the script exited early, was killed mid-flight, or simply set a restrictive mode without ever being re-flipped. We've stripped every `chmod` from the codebase. But files **written by older versions of those scripts** are still on disk in those restrictive modes (`0444` for `compiled_rules.txt`, `ip_blocklist.ipset`, `data/filters/<id>.txt`; `000` for `.env` after end-of-toggle lock). Re-running `toggle.sh` or `scripts/sync-rules.py` against those leftovers hits `PermissionError: [Errno 13] Permission denied` immediately.

Why the stale permissions also block `mv`/`rm` of the parent folder: The restrictive modes (`000` on `.env`, `0444` on output files) don't just block writes they prevent the kernel from unlinking the file's dentry during a `mv` or `rm` of the **containing directory** (`adguard/work/`). macOS enforces this at the VFS layer: you own the directory, but you can't delete a directory that contains entries you can't write to. This made the entire `adguard/work/` tree unmovable/undeletable until the locked inodes inside it were replaced. `unlock-files.sh` resolves this permanently by swapping each locked inode for a fresh `0644` one via atomic rename (`cp` + `mv`), which operates on the directory entry rather than the file contents no `chmod` called, no permission bypass needed.

````
**Problem:** You (the owner) cannot `cat`, `cp`, `rm`, `> redirect`, or any normal POSIX write against a `mode=000` file even though you own it  the basic mode bits block ALL access for owner, group, and other. The script does not call `chmod` and you want a permanent fix that never relies on a chmod from any future run either.

**Solution:** Replace the locked inode with a fresh readable one. `mv` is atomic; mode bits live in the inode, not the directory entry. Two flavors cover every case we hit on a real machine:

**Flavor A  locked file is mode `000` (owner can't even READ it):**
NOPASSWD sudoers (`/etc/sudoers.d/nullexit`) already includes `/usr/bin/python3`. So we read via `sudo python3` (which has the DAC override root has) and pipe the bytes around as base64 in user space, where the redirect happens with the user's normal umask = `0644`:
```bash
sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).decode())" > /tmp/snap.b64
base64 -d < /tmp/snap.b64 > <FILE>.new
mv -f <FILE>.new <FILE>
ls -la <FILE>          # mode is now 0644
```
The old `mode=000` inode is unlinked by `mv`; the new `mode=0644` inode (created by base64-decode via the calling user's umask) takes the path. No chmod ever called. The parent directory just needs to be writable by you (it always is, since you own it).

**Flavor B  locked file is mode `0444` (you can READ but cannot WRITE; `scripts/sync-rules.py` needs to rewrite):**
You own the file and can read it, you just cannot truncate/overwrite it. The simplest path is to rename it aside and let the writer create a fresh inode from scratch (which gets the umask-default `0644`):
```bash
mv <FILE> /tmp/old.<FILE>.$$
python3 scripts/sync-rules.py    # now writes <FILE> cleanly at mode 0644
```
Or apply Flavor A if you'd prefer to preserve the contents while resetting the mode.

**Why this is the canonical recovery, not a hack:**
- Mode bits live in the inode, not the path. `mv -f` replaces the path's inode reference; the old locked inode drops out of the directory entirely (dentry eviction) and loses its last link, so the kernel reclaims it. The new inode (whatever it is: a freshly-written one, or a base64-decoded copy) gets the path.
- The only prerequisite to apply Flavor B's `mv` is: you own the parent directory (so you can unlink entries from it) AND you have a NOPASSWD-capable binary that can read the byte stream if mode prohibits user read.
- This pattern generalizes. Any time a script ends up producing a "locked file that the next run can't update" situation, the same trick applies without a single chmod anywhere.

**Empirical verification (user machine, July 1, 2026):**
- `.env` was `mode=000` (owner $USER, i.e. you): Flavor A unblocked it in 3 commands, no chmod.
  - Before: `----------  1 $USER  staff  899 Jun 28 07:07 .env`
  - After:  `-rw-r--r--@ 1 $USER  staff  899 Jul  1 20:39 .env`
  - `docker compose config` immediately picked up `TS_AUTHKEY` + `WIREGUARD_PRIVATE_KEY` substitution.
- `compiled_rules.txt`, `ip_blocklist.ipset`, `data/filters/1782645604.txt` were `mode=0444`: Flavor B (`mv`-aside) unblocked all three.
  - scripts/sync-rules.py then ran clean: 279587 block rules + 171 allow rules + 16710 IP entries compiled; new files came out at `0644`.
- `toggle.sh` then went end-to-end in 48s (warp / routing-fix / adguardhome / tailscale all healthy, gateway mesh-joined, DNS hijack verified single-server).

**When this trick is appropriate vs. inappropriate:**
-  You own the parent directory and the file (typical desktop scenario).
-  The lock is a leftover from a removed `chmod` call that you want off disk forever.
-  NOPASSWD sudo includes at least one interpreter (`python3` on this system) that can be used to read mode-0 files. If it does not, add it first: `<USER> ALL=(ALL) NOPASSWD: /usr/bin/python3`.
-  The file is owned by another user (root, another account) and the lock is intentional  respect it; do not bypass another user's security boundary.
-  The parent directory is owned/writable only by another user  escalate properly via sudo, or stop and ask the user.
-  You do not actually own the file and there is a legitimate reason for its lock state  restore the file via a trusted source rather than circumventing the perms.

**Reusable cookbook:**

```bash
# Quick diagnostic: figure out which flavor you need.
WHOAMI=$(whoami)
ls -la <FILE>
stat -f 'mode=%Sp owner=%Su group=%Sg' <FILE>
# mode=---------- or mode=----r----- : you cannot even read it  Flavor A.
# mode=r--r--r-- but writer fails  : you are blocked from write but can read  Flavor B.
# mode=rw-r--r--                   : not actually locked at all; unrelated write failure.

# Flavor A (mode=000):
sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).decode())" > /tmp/snap.b64
base64 -d < /tmp/snap.b64 > <FILE>.new && mv -f <FILE>.new <FILE> && rm -f /tmp/snap.b64
````
````

```
# Flavor B (mode=0444, file is readable):
mv <FILE> /tmp/old.<FILE>.$(date +%s)
# (let the original writer recreate <FILE>; new file lands at mode 0644)

# Verify after either flavor:
ls -la <FILE>             # mode should be 0644
<your normal validation command>
```

15.9.7 sudo -n bash -c Silently Breaks the Scoped NOPASSWD Grant for dns-proxy.py (Fixed 2026-07-18)
Symptom: Found while verifying `scripts/pihole-mode.sh` actually works, not from a reported failure `bash scripts/pihole-mode.sh enable` hung on an interactive password prompt despite passwordless sudo being correctly configured and working for every other privileged call (`pfctl`, `networksetup`, etc. all succeeded fine under `sudo -n`).

Root cause: `sudo -l` shows the live NOPASSWD grant for the DNS proxy is scoped to the *literal* command `/usr/bin/python3 /Users/.../scripts/dns-proxy.py` (plus Homebrew-path variants) no shell wrapper, no env-var prefix. But `pihole-mode.sh`'s `cmd_enable()` invoked it as `sudo -n bash -c "LISTEN_ADDR=... TARGET_PORT=... python3 dns-proxy.py & echo $! > pidfile"`. sudoers matches the *literal argv* passed to `sudo`; here that's `bash -c "<string>"`, which doesn't match a rule scoped to `python3 <path>` at all so it silently fell through to requiring interactive auth. Even fixing just the wrapper wouldn't be enough: `dns-proxy.py` reads its config (`LISTEN_ADDR`, `LISTEN_PORT`, `TARGET_HOST`, `TARGET_PORT`) purely from environment variables, and `sudo`'s `env_reset` (macOS default) strips any custom env var not in `env_keep` none of these four are in the default `env_keep` list, and the sudoers rule carries no `SETENV` tag.

Same bug, dormant, in the primary codebase: `scripts/common.sh`'s `start_local_dns_proxy()` the SOCKS5-fallback DNS path used when exit-node pre-flight checks fail used the identical `sudo -n bash -c "TARGET_PORT=$DNS_PROXY_PORT ..."` pattern. Never caught because this session (like most) achieves exit-node mode every time, so the fallback path never actually runs. Had it ever been needed, the one thing meant to keep DNS working when the primary path fails would have hung on a password prompt with no TTY to answer it.

Fix: Rather than fight `sudo`'s env-passing restrictions (which would require editing the live sudoers file a change requiring a password I can't supply non-interactively, and not something to script unprompted), `dns-proxy.py` now optionally reads `/tmp/nullexit-dns-proxy.conf` (written `0600` by the caller, checked on top of the existing env-var/default values) if present. Both callers now write their overrides to that file, then invoke `sudo -n "$PYTHON" "$DNS_PROXY_SCRIPT"` directly the exact, unwrapped, already-whitelisted command, no `bash -c`, no env prefix. Both `stop_local_dns_proxy()` and `pihole-mode.sh disable` clean up the conf file. Verified end-to-end: `enable` no password prompt `test` resolves clean domains, sinkholes `doubleclick.net` `0.0.0.0` PASS `disable` removes both the process and the conf file.

Threat model note: the conf file is a bounded local-TOCTOU surface, same class already accepted for `/tmp/nullexit-ports.env` (15.9.5) restrictive perms, not a hard boundary against a co-resident attacker who could already do more damage with local code execution than redirect a DNS forwarder.

15.10 Tor Container

15.10.1 July 11, 2026: Tor Container Hardening & obfs4proxy Restoration
Symptom:
1. The Tor container entered a fatal crash loop upon boot, throwing `exec: "obfs4proxy": executable file not found in $PATH` when `TOR_USE_BRIDGES=true`.
2. When the user pasted Tor bridge lines containing comments (`#`) or blank lines into `tor-bridges.txt`, the container failed to validate the file properly and crashed.

Root Cause:
- **Bug 1 (Missing Pluggable Transports):** The Tor Dockerfile was built on `debian:bullseye-slim`, which had silently dropped or failed to resolve the `obfs4proxy` package in its latest apt repositories.
- **Bug 2 (Fragile Validation):** The `entrypoint.sh` bridge validation check merely checked `[ -s tor-bridges.txt ]` (file size > 0). It did not verify if the file actually contained valid, uncommented bridge lines. Passing empty lines or comments tricked the validation script into appending invalid `Bridge` directives to the `torrc`, causing the Tor daemon to instantly exit.

Resolution:
- **Upgraded Base Image:** Shifted the Tor Dockerfile base image to `debian:bookworm-slim`, restoring access to the `obfs4proxy` pluggable transport package.
- **Robust Bridge Validation:** Replaced the fragile `[ -s ]` check with a strict `grep -vE '^\s*(#|$)'` command that strips comments and empty lines. If no valid lines remain, the container safely falls back to standard public guard relays instead of crashing, ensuring Tor remains available even with a misconfigured bridge file.
- **Image Pinning:** Explicitly pinned `image: nullexit-tor:v1.0.0` in `docker-compose.yml` to prevent arbitrary local builds from drifting to `:latest`.

15.10.2 July 12, 2026: Tor ControlPort Dynamic Password Authentication & User Config
Symptom: Querying the Tor Control Port (e.g., via the `/sweep` script or `check_circuits.py`) returned empty responses or connection errors. Logs showed that Tor automatically closed its ControlPort:
`You have a ControlPort set to accept unauthenticated connections from a non-local address. ... That's so bad that I'm closing your ControlPort for you.`

Root Cause: Docker port mapping (especially under Colima on macOS) requires container sockets to bind to `0.0.0.0` to be reachable from the host. However, when Tor's ControlPort is bound to `0.0.0.0` (non-local)

```
     with no authentication, Tor closes it by default as a safety precaution. Sourcing SOCKS5/ControlPort to `127.0.0.1` inside the container was not an option as it broke the Docker bridge routing path.

**Resolution:**
- **Dynamic Password Authentication**: Updated the Tor container's `entrypoint.sh` to read `TOR_PASSWORD` and dynamically generate its HMAC hash via `tor --hash-password`, adding `HashedControlPassword <hash>` to `torrc`. This secures the ControlPort on `0.0.0.0` so Tor keeps it open.
- **Unified Credentials in `.env`**: Exposed `ADGUARD_USER=admin` and `ADGUARD_PASSWORD=nullexit` in `.env` to serve as the unified gateway credentials. Sourced `TOR_PASSWORD` directly from `ADGUARD_PASSWORD` in `docker-compose.yml`.
- **Client Authentication**: Updated `scripts/check_circuits.py` to fetch `TOR_PASSWORD` (defaulting to `ADGUARD_PASSWORD` or `nullexit`) and authenticate with the ControlPort using `AUTHENTICATE "<password>"` to regain access.

> See also 11.3 (Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification) for the design-level fail-closed guarantee.

### 15.11 macOS Platform Quirks

#### 15.11.1 macOS Application Firewall Silently Blocking AirDrop & Continuity
**Problem:** Users of complex network extensions (Tailscale, LuLu, Docker, etc.) often run into an issue where AirDrop, AirPlay, and Universal Clipboard silently fail, even with the gateway inactive and Wi-Fi/Bluetooth correctly configured. The internal Apple Wireless Direct Link (`awdl0`) interface appears healthy, but connections never establish.

**Root Cause:** The built-in macOS Application Firewall has a specific list of explicitly blocked applications. Apple's own core networking services (e.g., `/usr/libexec/sharingd`, `/usr/libexec/rapportd`, `/usr/libexec/avconferenced`) can be silently added to the "Block incoming connections" listeither through accidental "Deny" clicks on permission prompts, execution of aggressive privacy-hardening scripts, or a known macOS bug that retains strict blocks even after toggling the "Block all incoming connections" setting off. Since these are background daemons, macOS provides no visible error.

**Fix:** The firewall rules must be cleared via the terminal (or System Settings > Network > Firewall > Options):
```bash
sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/sharingd
sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/rapportd
```
Once unblocked, `sharingd` immediately resumes broadcasting mDNS/Bonjour discovery packets, and AirDrop restores instantly without requiring a reboot.

15.11.2 Standalone Tailscale Daemon Freeze after macOS Sleep/Wake
Problem: When the macOS host goes to sleep and wakes up, the network interfaces and routing tables are rebuilt. The standalone `tailscaled` daemon (installed via Homebrew) occasionally fails to handle this transition and becomes completely frozen/unresponsive. In this state, any command like `tailscale down` or `tailscale status` hangs indefinitely, and all DNS requests sent to the local Tailscale resolver (`100.100.21.8`) time out, causing a total internet blackout on the host.

Fix: Enhanced the Tailscale verification and teardown logic in `toggle.sh` and `recover.sh`:
1. **Unresponsive Daemon Detection:** Instead of assuming the daemon is healthy just because the Unix socket `/var/run/tailscaled.socket` exists, the scripts now actively run a status check with a 5-second timeout (`run_with_timeout 5 tailscale status`). If the check times out (exit code 143), the daemon is classified as unresponsive.
2. **Auto-Recovery Loop:** If the daemon is unresponsive, the script now automatically attempts to restart the service using `brew services restart tailscale` (falling back to `sudo -n brew services restart tailscale`).
3. **Graceful Retry:** Once restarted, the script pauses for 3 seconds for the daemon to initialize before proceeding with connection or teardown commands.

15.11.3 macOS Sharing Services (AirDrop/AirPlay) Freeze on Network Transitions
Problem: When the gateway is turned on or off, the scripts perform network configuration cleanups which include flushing the routing tables (`route -n flush`), restarting the `mDNSResponder` daemon (`killall -HUP mDNSResponder`), and power-cycling the Wi-Fi interface (`setairportpower` or `ifconfig en0 down/up`). While AirDrop BLE discovery continues to function, the actual file transfers stall at "Waiting..." indefinitely.

Root Cause: The rapid tear-down and rebuild of the local routing table and Wi-Fi interface causes Apple's core sharing and connection daemons (`sharingd` and `rapportd`) to lose their socket bindings to the mDNS/Bonjour discovery interface. Instead of reconnecting gracefully, they enter a wedged state and try to route peer-to-peer (AWDL) IPv6/TCP transfer traffic into the default gateway (the Tailscale interface `utun`) instead of the direct wireless link, causing the TLS verification handshake to time out.

Fix: Integrated an automatic sharing services reset into both `toggle.sh` and `recover.sh`:
1. The script now automatically runs `sudo -n killall sharingd rapportd` at the end of both the **START** path and the **STOP** path (as well as inside `recover.sh`).
2. Killing these daemons forces macOS to immediately relaunch them, binding them fresh to the newly initialized network interfaces and routing tables, allowing AirDrop to work seamlessly while the gateway is active.

15.11.4 Chrome Remote Desktop Connection Failures (Tailscale on Remote Device)
Context: When attempting to access a remote machine via Chrome Remote Desktop (CRD), the connection fails or hangs if Tailscale is active **on the remote device**. The connection works fine even if Tailscale (and
```

```
     the exit node) is active **on the local client/viewer device**, but only if Tailscale is disabled on the
     remote machine. *(This is distinct from the CRD-through-WARP latency limitation in 14.3.)*
1802
1803 **Why:** Having Tailscale active on the remote machine creates virtual network interfaces and DNS/routing
     modifications that conflict with Chrome Remote Desktop's WebRTC bindings and signaling traffic.
1804
1805 **Workaround:**
1806 - **Use Tailscale Native RDP/RustDesk (Recommended):** Connect directly to the remote device's Tailscale IP
     (`100.X.Y.Z`) via an RDP or RustDesk client. This is faster and avoids third-party relay conflicts.
1807 - **Disable Tailscale on the Remote Device:** If you must use CRD, disable Tailscale on the remote machine
     before connecting.
1808
1809 #### 15.11.5 tailscaled Daemon Broken State (brew services)
1810 **Symptom:** `brew services list` shows `started` but `tailscale status` shows "stopped".
1811
1812 **Fix:** `sudo brew services restart tailscale`. Toggle script now detects and auto-restarts.
1813
1814 #### 15.11.6 Apple Silently Removed the `airport` Binary in macOS Sonoma 14.4 (March 2024)
1815 **What happened:** `watcher.sh`'s `detect_lan_p2p_mode()` was written to use the `airport` CLI tool at:
1816 ```
1817 /System/Library/PrivateFrameworks/Apple80211.framework/Versions/Current/Resources/airport
1818 ```
1819 This path returned `exit 127: no such file or directory`. The function silently fell back to `reason="no-wifi"`
      even while the Mac was actively connected to WPA2 Enterprise Wi-Fi.
1820
1821 **Why Apple removed it:** `airport` was never a real public binary  it lived inside a **private framework** and
      was always technically unsupported. Apple deprecated it quietly years ago and finally removed it in **
     macOS Sonoma 14.4 (March 2024)**. The removal is part of a broader privacy clampdown on Wi-Fi metadata:
1822 - Starting in **macOS 14.5**, BSSIDs and other Wi-Fi association details are **redacted** (`<redacted>`) in
     most tools, including the official replacement `wdutil`
1823 - Apple's official recommendation is to use the **CoreWLAN framework** (Swift/Objective-C) with proper
     entitlements  useless for a bash script
1824 - The official CLI replacement `wdutil` requires `sudo` for most useful output  also useless for a background
     LaunchAgent
1825
1826 **Fix:** Replaced `airport -I | awk '/link auth/'` with:
1827 ```bash
1828 system_profiler SPAirPortDataType | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF
     ; exit}'
1829 ```
1830 This returns human-readable strings like `WPA2 Enterprise`, `WPA2 Personal`, or `None` for the currently
     associated network  without `sudo`, and without any removed private binary.
1831
1832 **Confirmed working output on this machine:**
1833 ```
1834 P2P detect: security='Enterprise'  allow=false (reason: 802.1x-enterprise)
1835 ```
1836
1837 > ** Risk: `system_profiler` may not survive forever either.** `system_profiler SPAirPortDataType` still works
     as of macOS Sequoia (15.x) without sudo and without redaction. However, given Apple's trajectory on Wi-Fi
     privacy, this could be locked down in a future release. If `security` comes back empty on a future macOS
     version while the Mac is actively on Wi-Fi, the function will silently fall back to `allow=false` (safe
     default), but the detection will be broken again. The long-term fix would be a small Swift helper using
     CoreWLAN that we compile once and bundle in the repo. *(This forward-looking caveat is also tracked as an
     observation in 16.)*
1838
1839 #### 15.11.7 Changing macOS Local Hostname
1840 **Context:** Users may wish to change their Mac's computer name/hostname in macOS System Settings.
1841 **Effect:** Changing the Mac's hostname will **not** break Nullexit (which relies on internal routing/localhost
     ). However, Tailscale will automatically detect this change and update the machine name on the Tailnet.
1842 - The underlying Tailscale `100.x.x.x` IPv4 address will remain exactly the same.
1843 - The **MagicDNS name** will change to match the new hostname. Users must remember to update their SFTP/SSH
     clients to use the new MagicDNS address.
1844
1845 #### 15.11.8 macOS MAC Address Spoofing (Apple Silicon)
1846 **Observation:** macOS prevents setting a new MAC address (`ifconfig en0 ether <mac>`) while the Wi-Fi
     interface is actively associated with an SSID. Additionally, some drivers on Apple Silicon completely
     reject spoofed MAC addresses via `ifconfig`, effectively ignoring the command and retaining the hardware
     MAC.
1847 **Workaround / Fix:**
1848 1. **Disassociate First:** The Wi-Fi interface must be turned off (`networksetup -setairportpower en0 off`)
     before attempting to set the MAC address, and turned back on afterward. This allows the OS to accept the
     change.
1849 2. **Fail-Safe Testing:** The `scripts/device-scramble.sh test-mac` command was implemented to safely verify if
      the underlying hardware and the current network will accept a randomized MAC address, as behavior varies
     significantly between Intel Macs (usually permissive) and Apple Silicon Macs (often restrictive).
1850
1851 #### 15.11.9 Shell-Language Landmines: macOS `/bin/bash` 3.2 + `set -e` (Field Notes)
1852
1853 This project's lifecycle scripts (`toggle.sh`, `recover.sh`, `common.sh`, ) run under `#!/bin/bash`, which on
     macOS is **GNU bash 3.2.57 (2007)**  Apple has never shipped a newer bash because bash 4+ is GPLv3. (
```

Homebrew's bash 5 lives at `/opt/homebrew/bin/bash` but is *not* the shebang target, and cannot be assumed on PATH a bare `bash --version` on a stock Mac reports 3.2.) Every script here also runs `set -e` (exit-on-error) with an `ERR`/`INT`/`TERM`/`HUP` trap. That combination **ancient bash + `set -e` + traps + `set -x` xtrace redirection** is a minefield. The traps below were all hit and fixed during development; document them so they are not re-discovered.

A. `set -e` + `if  then  else  fi` can abort at the `else` (bash 3.2 heisenbug).
The nastiest one. A construct as innocent as:
```bash
if [ -f "$MARKER_FILE" ]; then
  X="yes"
else
  X="no"
fi
```
aborted `toggle.sh` at init under `set -e` whenever the file was **absent** the false status of the `[ -f  ]` test leaked through the compound and `set -e` killed the script (`EXIT: code=1 phase=init`), before any real work ran. Two things made this vicious: (1) it is intermittent/context-dependent it did **not** reproduce in isolated `bash -c ''` snippets of the identical block; it only manifested inside the full script with the `set -x` xtrace fd-redirect (`exec 2> >(tee )`) active. We only proved causation by **bisecting the live script** (replace the block with a trivial assignment the script sailed past init). (2) Static tools (`bash -n`, `shellcheck`) cannot see it it is a runtime interaction. **Safe pattern:** default-assign, then an `if` with **no `else`** (an `if` whose false condition has no else-branch yields status 0):
```bash
X="no"
if [ -f "$MARKER_FILE" ]; then X="yes"; fi
```
Confirmed root cause (2026-07-13, reproduced head-to-head). The trigger is specifically **`set -x` plus a `PS4` that contains a command substitution** here the `DEBUG_TRACE` prompt `PS4='+ [$(date )] '`. Forcing the false-condition/else path (`gateway_state` `stopped`, so `if [ "$(gateway_state)" = "running" ]` takes the else/START branch) on stock `/bin/bash` 3.2.57: with the **default `PS4`** it survives (3/3, exit 0); with toggle's **`$(date)` `PS4`** the `ERR` trap fires on the else branch (`rc=1` from the false condition, 3/3, exit 42). The `$(date)` runs a subshell on *every* traced command; while tracing the `if`-condition it perturbs `$?` so the condition's non-zero status leaks past the normal `if`-exemption and trips the `ERR` trap. This is a true **heisenbug** enabling tracing (`DEBUG_TRACE=true`) is what *causes* the abort, which is exactly why it's invisible with `DEBUG_TRACE=false` (the default) and why every isolated repro on the *default* `PS4` passes (this misled an entire debugging session into briefly "disproving" the landmine). It aborted the gateway START decision (`toggle.sh:689`) under `DEBUG_TRACE=true`; with `DEBUG_TRACE=false` the gateway starts cleanly (verified: 192 s, all five containers healthy). **Fixed (2026-07-14).** `DEBUG_TRACE`'s `PS4` now uses the subshell-free `${SECONDS}` (elapsed seconds since the shell started) instead of `$(date "+%H:%M:%S")`, removing the per-traced-command subshell that perturbed `$?`. bash 3.2 has no subshell-free *wall-clock* token for `PS4` (`$EPOCHSECONDS` is 5.0+, `printf '%()T'` is 4.2+), so the per-line stamp is now elapsed-seconds the absolute start time is still logged once in the `DEBUG_TRACE  enabled ` header line, so wall-clock is recoverable. Bonus: this also drops a `date` fork on *every* traced line (previously thousands). **Verified head-to-head on stock `/bin/bash` 3.2.57:** `toggle.sh --restart` with `DEBUG_TRACE=true` now completes `code=0` with tracing active straight through the exact decision that aborted before the trace shows `gateway_state  stopped` (`common.sh:229`) feeding the false `if [ "$(gateway_state)" = "running" ]`, then the else branch reaching `Action: STARTUP GATEWAY` (1436 traced lines, no `ERR` trap firing, all five containers healthy). `DEBUG_TRACE` is now safe to use on the STOP/START decision path.

More generally, any command whose "failure" is normal control flow (`grep -q`, `[  ]`, `diff`, `kill -0`) that ends up as the *last* command of a function/branch/script under `set -e` is a landmine; guard it (`|| true`) or restructure so a non-zero status can't be the tail value.

B. `BASH_XTRACEFD` and the `exec {var}>>file` dynamic-fd form are bash 4.1+ (hard-fail on 3.2).
Our `DEBUG_TRACE` feature wants xtrace (`set -x`) sent to `output.log` while keeping the terminal clean the clean way is `exec {BASH_XTRACEFD}>>"$LOG_FILE"`. On 3.2 this is a **syntax/runtime error**: `exec: {BASH_XTRACEFD}: not found` (dynamic-fd allocation `{var}>>` didn't exist until 4.1, and `BASH_XTRACEFD` itself is 4.1+; on 3.2 it's just an ordinary variable that `set -x` ignores). Left unguarded it would abort the script at startup i.e. the debug feature would break the very thing it's meant to observe. **Safe pattern:** branch on `${BASH_VERSINFO[0]}.${BASH_VERSINFO[1]}` on 4.1 use `exec 9>>"$LOG_FILE"; export BASH_XTRACEFD=9`; on 3.2 fall back to `exec 2> >(tee -a "$LOG_FILE" >&2)` (xtrace goes to stderr, which we duplicate to the log; the terminal also shows it, which is acceptable). Never use the `exec {var}>>` form unless you *know* you're on 4.1.

C. `trap  ERR` + `$LINENO` misattributes the failing line on 3.2.
When a command fails under `set -e`, our trap runs `cleanup_handler ERR $LINENO`. On bash 3.2, `$LINENO` inside an `ERR` trap frequently reports the line of the **enclosing `fi`/closing brace of a compound statement**, not the true failing line. During the STOP/START-decision bug this produced `Script failed at line 1202` where 1202 is a bare `fi` actively misleading. **Mitigation:** don't trust the trap's line number as gospel on macOS; corroborate with the `EXEC:`/`EXIT:` lifecycle breadcrumbs and the `DEBUG_TRACE` xtrace (which prints `file:line:function` via a custom `PS4`), which pin the *actual* last-executed command.

D. `set -e` is not inherited by command substitutions / subshells the way people expect, and is suppressed inside `if`/`&&`/`||` conditions.
This cuts both ways and is a frequent source of "why didn't it fail?" and "why *did* it fail?": a non-zero command used as an `if` condition (or left/right of `&&`/`||`) is *exempt* from `set -e` (by design), but the moment that same call is a bare statement it aborts. `is_gateway_active()` relies on this (it's always called as `if is_gateway_active; then `), which is precisely why moving the decision to a marker file

instead of a probe that must be `if`-guarded to be safe is more robust (see 15.12.19). Also note `local X=$(cmd)` masks `cmd`'s exit status (the `local` builtin's status wins) a `shellcheck` SC2155 we see repeatedly; harmless when we don't check `$?`, but a real footgun if you do.

E. Process substitution (`>()`, `<()`) is a bashism, and its lifecycle is loose.
Our 3.2 xtrace fallback (`exec 2> >(tee -a "$LOG_FILE" >&2)`) uses process substitution fine under bash, but (1) it is **not** POSIX `sh`, so these scripts must never be invoked as `sh script.sh`; and (2) the background `tee` it spawns is reaped asynchronously, so the very last few lines written just before exit can occasionally be truncated in the log. Acceptable for a debug trace; do not rely on it for critical last-gasp logging (use a direct `>> "$LOG_FILE"` append for must-capture lines, as the `EXIT` breadcrumb does).

F. GNU-vs-BSD userland divergence (not bash itself, but same class of pain).
macOS ships BSD variants of the core tools, so idioms differ from Linux: `sed -i ''` (BSD requires the empty backup-suffix arg) vs `sed -i` (GNU); `stat -f%z` (BSD) vs `stat -c%s` (GNU); `date` format flags; `jot` (BSD) vs `shuf` (GNU) for random port generation; `route`/`networksetup`/`dscacheutil` (macOS-only) vs `ip`/`resolvectl` (Linux). This is *why* the codebase branches on `[[ "$OSTYPE" == "darwin"* ]]` everywhere `toggle.sh` and `recover.sh` are each single cross-platform scripts that dispatch on `$OSTYPE` (an early Linux block runs and exits before the macOS body). `timeout(1)` doesn't exist on stock macOS at all hence the pure-bash `run_with_timeout` in `common.sh`.

Working rules distilled from the above:
- Assume **bash 3.2** semantics for anything under `#!/bin/bash` on macOS; gate any 4.0+ feature (`BASH_XTRACEFD`, dynamic `{var}` fds, associative arrays, `${var^^}` case-conversion, `mapfile`/`readarray`, negative array indices) behind a `BASH_VERSINFO` check or avoid it.
- Under `set -e`, never let a "benign non-zero" command (`[ ]`, `grep -q`, `kill -0`, `diff`) be the tail statement of a branch/function/script; use no-`else` `if`s, `|| true`, or explicit `$?` handling.
- Prefer **persisted state (marker files)** over **live probes** for control-flow decisions a probe forces an `if`-guard and can hang/abort; a file read cannot (see 15.12.19).
- `bash -n` and `shellcheck` catch syntax and many smells but **cannot** catch `set -e`/trap runtime interactions a **live run is mandatory** to validate lifecycle-script changes; the `DEBUG_TRACE` + breadcrumb instrumentation exists precisely to make those live runs diagnosable.
- Never run these scripts as `sh` they are bashisms end to end.

15.11.10 The Live Wi-Fi SSID Is Unreadable From Any Script (Darwin 25) Confirmed Including Root
Finding (2026-07-23, while building Wi-Fi auto-rejoin, 15.5.9). On this Mac (Darwin 25 / macOS 26) there is **no** script- or root-accessible path to the *currently associated* SSID. This was verified empirically against every candidate:

- `networksetup -getairportnetwork en0` "You are not associated with an AirPort network" *even while connected*.
- `ipconfig getsummary en0`, `system_profiler SPAirPortDataType`, `wdutil info`, `scutil` the SSID field prints `<redacted>`. `wdutil` redacts **as root too**. The legacy `airport` binary was removed back in 14.4 (15.11.6).
- The on-disk store `/Library/Preferences/com.apple.wifi.known-networks.plist` (per-network SSIDs + timestamps) is `0600 root:wheel` with **no** `restricted` flag, yet `sudo python3 open()` still returns `PermissionError` it lives in a macOS **data vault** (same protection class as `TCC.db`), which denies even root regardless of flags. `defaults`/`plutil` can't reach it either. `airport.preferences.plist` is SIP-protected.

The **only** dynamic path to the live SSID is a **Location-Services-authorized** process (`shortcuts` CLI, or a CoreWLAN+CoreLocation helper). The user explicitly refused to grant Location, so the live SSID is treated as permanently unobtainable. This is *why* Wi-Fi auto-rejoin (15.5.9) is built on a user-**remembered** name list rather than reading the current network joining *by name* is the one operation Apple leaves ungated. Note this is distinct from 15.11.6, which is specifically about the `airport`-binary removal breaking `detect_lan_p2p_mode`'s security probe; `system_profiler`'s `Security:` field still works, only the SSID itself is redacted.

15.12 Misc Resolved

15.12.1 Gluetun Resets nftables FORWARD Rules
Symptom: FORWARD RELATED,ESTABLISHED rule disappears after Gluetun health check.

Fix: `scripts/routing-fix.sh` re-adds to both iptables backends every 30 seconds. Legacy backend (policy ACCEPT) acts as safety net.

15.12.2 Native SSH & SFTP Security (Zero Local Attack Surface)
Context: macOS "Remote Login" (`sshd`) and "File Sharing" (`smbd`) open ports on the local network (e.g. `172.x.x.x`), exposing the host to brute-force attacks on public Wi-Fi.

Implementation: The script strictly enforces `tailscale up --ssh=true` whenever the mesh state is reset. This empowers the user to completely disable native macOS Remote Login and File Sharing. This provides an incredible zero-trust security advantage: even if an attacker steals your Mac username and password, they **cannot** access your files remotely. Disabling the native services completely closes the listening ports on the local Wi-Fi interface (`172.x.x.x`), rendering the Mac inaccessible to local network logins. Meanwhile, Tailscale intercepts port 22 traffic exclusively over the encrypted mesh and authenticates cryptographically based on your Tailscale SSO identity. Stolen Mac passwords are mathematically useless without first bypassing your Tailscale Two-Factor Authentication and registering a device onto the mesh.

Cross-Platform File Exchange: To easily browse and exchange files securely, simply use an SFTP client and connect to your Mac's MagicDNS name on port 22. Recommended clients: **WinSCP** or **FileZilla** (Windows), **Cyberduck** (macOS), and **FE File Explorer** or **Solid Explorer** (iOS/Android).

```
1917
1918 #### 15.12.3 Logging Architecture
1919 For debugging, logs are strictly segmented based on the component's lifecycle:
1920 - **`output.log` (Host-side):** Contains all standard error (`stderr`) and verbose output from the host scripts
       (`toggle.sh`, `recover.sh`, `setup.sh`). Since the `rule-compiler` container is ephemeral and deleted
      after running, its logs (and any Python errors from `scripts/sync-rules.py`) are extracted and appended to
      this file before deletion.
1921   - **Log Rotation Policy:** On startup, `toggle.sh` checks if `output.log` exceeds **1GB** (`1,073,741,824`
      bytes). If it does, it performs a metadata-only rename (`mv output.log output.log.old`), preserving history
       while resetting the active log to 0 bytes. This rename-based rotation avoids the heavy disk I/O and CPU
      overhead of rewriting a massive file line-by-line. The threshold is intentionally large because with `
      DEBUG_TRACE=true` the log accumulates a full command-by-command execution trace (effectively the framework'
      s entire compilation/warning history), which is valuable to retain for post-mortems; a smaller cap would
      rotate that history away too quickly.
1922   - **Egress Probe Logging:** The background host-egress leak prober checks if stdout is a TTY (`[ -t 1 ]`) and
       will only print live status updates (using carriage return `\r`) or duplicate console warnings in
      interactive mode. It automatically detects macOS/BSD vs. Linux environments to dynamically construct
      timestamps with date and time (omitting milliseconds on macOS to prevent high CPU utilization from spawning
       sub-second processes).
1923 - **`docker logs <container>` (Guest-side):** The persistent containers (`warp`, `tailscale`, `routing-fix`, `
      adguardhome`) use the Docker `json-file` logging driver with a strict `max-size` (1m-10m) to prevent VM
      disk exhaustion. Use standard `docker logs` to view them.
1924
1925 #### 15.12.4 Pre-Flight Check 1: Wrong `tailscale ping` Flag
1926 **Problem:** Check 1 of the pre-flight connectivity checks used `--c 1` (double dash) but `tailscale ping`
      accepts `-c 1` (single dash). The invalid flag caused the command to exit with code 1, marking the check as
       FAIL even though the gateway was reachable.
1927
1928 **Fix:** Changed `--c 1` to `-c 1`.
1929
1930 #### 15.12.5 Pre-Flight Check 1: DERP Relay Exit Code
1931 **Problem:** Even with the correct `-c 1` flag, `tailscale ping` exits with code 1 when the connection goes
      through a DERP relay (`direct connection not established`)  even though a pong was successfully received
      (28ms via DERP(tor)). The check treated relayed connections as failures.
1932
1933 **Fix:** Changed the check from relying on exit code to piping stdout through `grep -q "pong"`. This accepts
      relayed connections as valid (they work fine for the exit node use case), while still failing on true
      timeouts or unreachable peers.
1934
1935 #### 15.12.6 IPv6 Exit-Node Leak
1936 **Problem:** Tailscale supports IPv6, but WARP/gluetun is IPv4-only. IPv6 packets forwarded through the exit
      node would leak through the Docker bridge unencrypted.
1937
1938 **Fix:** Added `ip6tables -A FORWARD -i tailscale0 -j DROP` to `post-rules.txt`, blocking all IPv6 forwarded
      traffic from the Tailscale interface. *(This is the fix referenced by the IPv6-leak scenario in 15.6.1.)*
1939
1940 #### 15.12.7 SOCKS5 Proxy Threading Race Condition
1941 **Problem:** The `forward` function in `scripts/socks5-proxy.py` called `select.select([src, dst], [], [])`
      which raised `ValueError: file descriptor cannot be a negative integer (-1)` when one thread closed a
      socket while the other thread was selecting on it.
1942
1943 **Fix:** Added `ValueError` exception handling around `select.select()` and `recv()` calls. When a file
      descriptor becomes negative, the thread breaks out of the forwarding loop cleanly instead of crashing.
1944
1945 #### 15.12.8 Docker Healthcheck: Multi-line Python in CMD Array
1946 **Problem:** The socks-proxy healthcheck used a `CMD` array with a multi-line Python string. YAML collapsed the
       newlines, causing the `#` comment to comment out the rest of the code and the `assert` statement to be
      mangled into `as`.
1947
1948 **Fix:** Changed from `["CMD", "python3", "-c", "..."]` to `["CMD-SHELL", "python3 -c '...'"]` with a single-
      line Python expression. Uses a real SOCKS5 handshake (sends `[5, 1, 0]`, expects `[5, 0]`) instead of a
      simple TCP port check.
1949
1950 #### 15.12.9 Graceful Shutdowns Leave DNS Hijacked
1951 **Problem:** The gateway overrides the macOS host's DNS settings (via `networksetup`) to route queries to the
      AdGuard container. Because macOS persists `networksetup` changes to disk, shutting down the Mac while the
      gateway is running causes the Mac to boot up later with hijacked DNS. Since the Colima VM and Docker
      containers don't auto-start on boot, the user would have no internet access until they manually run the
      toggle script.
1952
1953 **Root Cause:** The `toggle.sh` script exits after the gateway starts. There was no background daemon listening
       for macOS system shutdown signals to revert the network settings.
1954
1955 **Fix:** Wrapped the background `caffeinate` process (used to prevent system sleep) in a bash `trap` command.
      The trap listens for `SIGTERM`, `SIGINT`, and `SIGHUP`. When the user initiates a graceful shutdown, macOS
      sends `SIGTERM` to the background trap, which instantly executes `recover.sh` to flush the host DNS back to
       `1.1.1.1` right before the machine powers off.
1956 *Note:* We explicitly use `recover.sh` instead of `toggle.sh` because during a shutdown, the OS limits process
      cleanup time and is already independently sending termination signals to Docker and Colima. Trying to run `
      docker compose down` inside a shutdown trap creates a race condition that can hang the script until macOS
      forcefully `SIGKILL`s it. `recover.sh` abandons container management and surgically flushes the macOS
```

```
    network settings in milliseconds, guaranteeing completion before the OS pulls the plug.
1957
1958 #### 15.12.10 Linux Native Wi-Fi Bouncing
1959 **Problem:** The generated `scripts/recover-linux.sh` incorrectly retained the macOS `networksetup` commands,
         which would crash and fail to bounce Wi-Fi on Linux hosts.
1960
1961 **Fix:** Swapped the macOS logic for native Linux commands. The script now attempts `nmcli radio wifi off/on`
         first, and gracefully falls back to `ip link set wl... down/up` if NetworkManager is unavailable.
1962
1963 #### 15.12.11 TS_AUTH_ONCE Cloud Footgun & Ephemeral Keys
1964 **Decision:** The compose file uses `TS_AUTH_ONCE=true` to prevent authentication loops. However, on headless
         cloud VPS deployments, if a standard auth key expires (usually 90 days), the container will silently drop
         off the mesh. We documented this trade-off explicitly in `docker-compose.yml` and recommended the use of **
         Tailscale Ephemeral Keys** for cloud deployments, which automatically clean themselves up and bypass this
         issue.
1965
1966 #### 15.12.12 Hardcoded AdGuard Credentials
1967 **Decision:** AdGuard is deployed with the hardcoded credentials `admin / nullexit`. This is perfectly safe
         behind a NAT on a local macOS laptop. However, if deployed on a cloud host with exposed ports, this becomes
          a severe security risk. We added a stark warning comment in `docker-compose.yml` to ensure cloud users
         change this before deployment.
1968
1969 #### 15.12.13 Routing Fix: 30-second polling vs Netlink Socket
1970 **Decision:** Instead of building a complex Netlink socket listener (`ip monitor`) to react instantly to
         iptables and routing flushes from Gluetun, we retained the 5-second `sleep` loop in `scripts/routing-fix.sh
         `.
1971 - **Reasoning:** Building a netlink listener in pure bash on Alpine is incredibly brittle and complex (
         especially for iptables events). The 30-second polling loop is bulletproof. The only trade-off is a
         potential 1-4 second stutter if Gluetun reconnects, which was deemed an acceptable edge case for absolute
         reliability.
1972
1973 #### 15.12.14 Cache Poisoning and URL Rot in scripts/sync-rules.py
1974 **Problem:** If a remote blocklist URL went permanently offline (404 Not Found), the Python `urllib` request
         would return the 404 HTML string. The script would aggressively regex-parse this HTML, find no domains, and
          overwrite the healthy local cache with a 0-domain file. This effectively disabled ad-blocking until the
         URL was fixed.
1975
1976 **Fix:** Implemented a defensive programming sanity check in `scripts/sync-rules.py`. Before overwriting the
         cache, the script now verifies that the compiled domain count is greater than 10 (and 1 for IP feeds). If
         it falls below this threshold (indicating a catastrophic 404 failure across multiple URLs), it aborts the
         cache overwrite, raises a `ValueError`, and keeps the previous day's healthy cache intact.
1977
1978 #### 15.12.15 Gluetun nf_tables Parser Crash (Bash Interpolation Failure)
1979 **Problem:** Attempting to make the TCP MSS dynamically configurable by using a standard bash variable inside `
         post-rules.txt` (`iptables ... --set-mss ${GATEWAY_MSS:-1120}`) caused the Gluetun container to
         catastrophically crash during startup. Gluetun's internal parser executes `post-rules.txt` line by line
         without a shell, meaning the variable was never expanded. It literally attempted to inject the string `${
         GATEWAY_MSS:-1120}` into iptables, resulting in a fatal `bad value for option "--set-mss"` error.
1980
1981 **Fix:** Removed the bash variable from `post-rules.txt` and reverted it to a pure hardcoded integer. To retain
          dynamic `.env` configuration, we moved the interpolation logic to `toggle.sh`. Right before Docker starts,
          the host script parses `GATEWAY_MSS` from the `.env` file and uses a cross-platform `sed` command to
         dynamically overwrite the integer directly inside `post-rules.txt`. This perfectly mimics manual
         configuration and completely bypasses Gluetun's strict parser.
1982
1983 #### 15.12.16 AdGuard Filter Redundancy & Memory Optimization
1984 **Problem:** AdGuard Home does not perform cross-list deduplication in memory. When it loads its own native
         subscription filters (e.g., `AdGuard DNS filter`) alongside our massive `compiled_rules.txt`, overlapping
         rules are loaded twice, wasting significant RAM in the Colima VM. The UI would show ~500k rules,
         representing the sum of all lists rather than unique domains.
1985
1986 **Fix:** Updated `scripts/sync-rules.py` to intelligently cross-reference AdGuard's configuration and
         deduplicate our list against it.
1987 1. The script now reads `adguard/conf/AdGuardHome.yaml` to dynamically fetch the exact URLs of any enabled
         native AdGuard filters.
1988 2. The parsing engine was upgraded to normalize basic AdGuard syntax (`||domain^`) back into raw base domains
         during the build process.
1989 3. The script subtracts the native AdGuard domains from our custom blocklist *before* compiling, immediately
         purging ~83,000 completely redundant rules.
1990 4. The normalized base domains then pass through our subdomain optimizer, which squashes an additional ~50% of
         the remaining rules.
1991
1992 This reduces the final compiled output from ~325k to ~281k rules on the `heavy` profile, saving ~45k rules from
          being loaded into memory twice and reducing the overall RAM footprint of the `adguardhome` container.
1993
1994 #### 15.12.17 Kernel-Level IP Blocklist (Threat Intelligence Firewall)
1995 **Problem:** DNS sinkholing cannot stop malware or botnets that bypass DNS entirely by hardcoding direct IP
         addresses (a common technique for C2 communication). These connections pass straight through AdGuard unseen
         .
1996
```

Fix: Added a second compilation pipeline to `scripts/sync-rules.py` that runs on every startup alongside the DNS pipeline.
1. The compiler concurrently fetches four curated threat-intelligence feeds: **Feodo Tracker** (abuse.ch botnet C2 IPs), **Spamhaus DROP** (IPs allocated to criminal organizations), **Emerging Threats** (Proofpoint active C2 and scanners), and **CINS** (active brute-force sources).
2. All entries are normalized via Python's `ipaddress` module. Any IP or CIDR that overlaps with RFC1918 private ranges, loopback, link-local, or the Tailscale CGNAT range (`100.64.0.0/10`) is stripped to prevent accidentally locking users out of their own LAN or mesh.
3. The cleaned list (16,721 unique IPs/CIDRs) is written as an `ipset restore` file using an **atomic swap pattern**: `create_new  populate  swap with live  destroy_new`. This guarantees the live `block_malicious` ipset is never empty during a reload.
4. `scripts/routing-fix.sh` watches the file's mtime every 30 seconds. On change it runs `ipset restore` and idempotently re-injects `FORWARD DROP` rules for both `src` and `dst` into both iptables backends (nftables + legacy).
5. `docker-compose.yml` mounts `adguard/work/userfilters/` into `routing-fix` as `/userfilters:ro` and adds `rule-compiler: service_completed_successfully` to its `depends_on`, eliminating the race condition where routing-fix could start before the file existed.

Memory cost: The entire 16,721-entry ipset costs only ~1.6 MiB of kernel memory negligible in the 600MB VM budget.

15.12.18 Incident Post-Mortem: Discarded Docker Exec Errors in logs (July 10, 2026)
Symptom: When the `warp` container is stopped, dead, or failing to connect to its endpoints, the `toggle.sh` and `recover.sh` connectivity health checks fail, but no details are printed to `output.log`. The logs only show generic `FAIL` outputs, making it hard to see if the failure was a network timeout, Docker daemon error, or a missing binary.

Root Cause: Overly broad error silencing. The `docker compose exec` commands in the health checks (Check 3 in `toggle.sh` and the traffic check in `recover.sh`) both redirected stderr to `/dev/null` (`&>/dev/null` and `2>/dev/null` respectively), completely throwing away any diagnostic error messages produced by Docker or `wget`.

Fix (July 10, 2026): Redirected stderr of the `docker compose exec` check commands to `output.log` (`2>> output.log`). This ensures that any command execution failures or connection timeout details are logged.

15.12.19 Incident Post-Mortem: Toggle Decided STOP/START From a Live Probe, Aborting When Docker Was Down (July 13, 2026)
Symptom: Running `./toggle.sh` (or `--restart`) while the Colima VM / Docker daemon was **not** running caused the script to hang for ~15 seconds and then abort during the START phase without ever booting Colima. With `DEBUG_TRACE` the lifecycle breadcrumb showed `toggle.sh EXIT: code=1 phase=start`, and the trace showed execution jumping straight from the START-branch entry to the `ERR` trap (`cleanup_handler ERR`) with the START body never executing. Historically this silent failure occurred 57 in `output.log`. Spam-clicking the toggle during the resulting window (each retry rejected by the lock guard, then finally succeeding) is what produced the earlier "pile of stacked toggle processes" and Wi-Fi-bounce confusion.

Root Cause: `toggle.sh` decided whether to STOP or START by calling `is_gateway_active()` (`scripts/common.sh`), which **live-probes** the running system its first step is `run_with_timeout 15 docker compose ps --status running`. When `/var/run/docker.sock` is absent (Colima down), that call blocks for the full 15 s timeout and returns non-zero. Combined with the script's global `set -e` and `ERR` trap, a non-zero result evaluated at the `if is_gateway_active; then  else  fi` boundary aborted the whole script *before* the `else` (START) branch body ran precisely the branch whose job is to start Colima. Design-wise the deeper flaw is that a **disable** should not depend on whether the gateway is *currently* healthy; it should act on whether the gateway is *supposed* to be running (persisted intent). The correct signal already existed `MARKER_FILE=/tmp/nullexit-gateway-active.marker`, written at the end of a successful START and cleared on STOP but `toggle.sh` only wrote/cleared it and never read it for this decision. **Note:** this was a long-standing latent bug, not a regression from the July 2026 `common.sh` platform-function unification or the `DEBUG_TRACE` work; those changes only made the previously-silent abort *visible* (via the `EXEC:`/`EXIT:` breadcrumbs and xtrace).

Fix (July 13, 2026): Introduced `gateway_should_be_running()` in `common.sh` and routed all four decision sites through it (`toggle.sh` main decision + `--restart` pre-check; `scripts/toggle-linux.sh` main decision + `--restart` pre-check). It decides from persisted intent:
1. `toggle.sh` captures the marker's existence into `GATEWAY_MARKER_AT_STARTUP` **before** the top-of-script "defensive stale-marker clear" wipes it, and the helper honors that captured value (falling back to the live marker file for callers that run before the clear, e.g. the `--restart` pre-check).
2. Marker present STOP path; marker absent START path. This is instant and cannot hang or abort.
3. Only when the marker is absent **and** the Docker socket is actually reachable (`/var/run/docker.sock` or `~/.colima/default/docker.sock` on macOS) does it consult `is_gateway_active` as a non-fatal reconciliation fallback (covers a marker lost to a `/tmp` wipe while containers are genuinely up). A dead socket short-circuits to "stopped" with no probe, so `set -e` can never be tripped.
4. The STOP path's `docker compose down` is now `|| true`-guarded so a disable always completes even with Docker down; the START path's `docker compose up` is deliberately left unguarded (a failed start *should* trigger cleanup). On Linux the marker isn't maintained, so the helper degrades to the `is_gateway_active` fallback fast there because native Docker has no Colima VM to hang.

How we got here (methodology note). This one is worth recording because the *process* mattered as much as the patch. The failure had been happening silently for a long time the log simply ended after a teardown with no error, so there was nothing to chase. We first attacked the **observability gap** rather than guessing at the cause: we wrapped the heavyweight lifecycle commands (`colima start`, `docker compose up/down`) so their output and exit codes always land in `output.log`, added an `EXIT: code= phase=` breadcrumb that fires even when the process is killed, and put a full opt-in xtrace behind `DEBUG_TRACE`. Only *then*

did we reproduce the failure and this time the trace said, unambiguously, `EXIT: code=1 phase=start` with the START body never entered. The instrumentation didn't fix anything, but it converted an invisible abort into a precise, reproducible fact. (Lesson, consistent with 15.12.18 and the LUKS post-mortem referenced in 15.11: **a mechanism that fails silently is worse than one that fails loudly** so we make it fail loudly first.)

The root-cause leap, though, came from a **design-smell intuition, not the stack trace**: the observation that it is nonsensical for a *disable* to first check whether the gateway is *currently online* a teardown should act on whether the gateway is *supposed* to be running, not probe a subsystem to find out. That instinct turned out to name the bug exactly. The smell was a **conflated responsibility with an inverted dependency**: to decide whether to *start* the gateway, the code first had to *talk to* the gateway (via Docker) the very thing START is responsible for bringing up so the check was guaranteed to be unavailable in precisely the cold-start case that mattered. A second tell was **false symmetry**: enable (builds state) and disable (tears it down) are not symmetric operations, yet both ran the identical expensive probe; whenever two operations that should differ share suspiciously identical machinery, the code is usually lying about the structure of the problem, and reality eventually calls the bluff. The fix invented nothing the honest signal (`MARKER_FILE`, literally "the gateway is supposed to be up") already existed and was maintained in all the right places; it was simply never consulted for the decision it was made for. We just pointed the decision at the answer the codebase already knew. General principle worth keeping: **design smells here are load-bearing they tend to *be* latent bugs, not merely cosmetic**, and are worth chasing on intuition even before a trace confirms them.

Acceptance verified (static + live): with marker present STOP (instant, even with Docker down); with marker absent + Docker down START (0 s, no `set -e` abort); a real run now proceeds past init to the gateway-status decision. `bash -n` clean on all three scripts; no new `shellcheck` findings; re-signed `.signatures` (`toggle.sh`/`common.sh`/`toggle-linux.sh` are in the signed set) and `crypto.sh --verify` exits 0.

Follow-up (same day) a `set -e` regression the fix introduced, and why only a live run caught it. The first cut of the marker capture wrote the intent as a full `if [ -f "$MARKER_FILE" ]; then GATEWAY_MARKER_AT_STARTUP="yes"; else GATEWAY_MARKER_AT_STARTUP="no"; fi`. On macOS `/bin/bash` 3.2, with `set -e` active and the `DEBUG_TRACE` xtrace redirect (`exec 2> >(tee )`) in effect, that construct **aborted the script at init** the moment the marker was absent the false `[ -f ]` status leaked through the compound and `set -e` killed the run (`EXIT: code=1 phase=init`), before the gateway logic ran at all. Notably this did **not** reproduce in isolated `bash -c` snippets of the same block it only surfaced in the script's real runtime context so we had to bisect against the actual `toggle.sh` (swap the block for a trivial assignment the script sailed past init) to prove causation. Fix: write it as a default assignment plus an `if` **without an `else`** (`GATEWAY_MARKER_AT_STARTUP="no"` then `if [ -f "$MARKER_FILE" ]; then GATEWAY_MARKER_AT_STARTUP="yes"; fi`), which yields status 0 when the condition is false. Lesson reinforced: `set -e` on legacy bash is genuinely treacherous around conditionals, and the observability work paid for itself again the breadcrumb pinned the abort to init instantly, and a live run (not static checks) was the only thing that exposed a heisenbug invisible to isolated reproduction.

15.12.20 Honey-Port Tripwire Accumulation (July 14, 2026)
Symptom: After many toggle-ons, `ps` showed a pile of idle `bash ./toggle.sh` subshells, each blocked on an `nc -l localhost <port>` listener one leaked per toggle-on, never reaped.

Root Cause: The Honey-Port Tripwire arms a fresh random port each START via a disowned `( nc -l  ) &` subshell. `nc -l` blocks until a connection that (by design) almost never comes, so each toggle-on leaks one idle listener; because the port is random, successive listeners never collide and simply stack up over time.

Fix: `toggle.sh` now reaps the previous tripwire before arming a new one `pkill -f "nc -l localhost"` (macOS) / `pkill -f "nc -l .*127.0.0.1"` + `"nc -l localhost"` (Linux) immediately before `HONEY_PORT=$(get_free_port)` at both arming sites. No `sudo` needed (the listeners are user-owned). Net effect: at most one tripwire alive at a time instead of one per toggle. (The same `pkill` pattern cleared the backlog that this session's ~15 test toggles produced.)

15.12.21 Post-Wake Self-Feedback Loop: The Watcher Re-Fires On The Route Change It Caused (July 14, 2026)
Symptom: After a container restart (or wake), `output.log` showed `recover.sh --post-wake` launching ~6 times in a row, one launch roughly every ~40s, each exiting `exit=0` recovery *worked* every time yet kept re-firing "until it worked." Every launch was triggered by `NET:  changedKey State:/Network/Global/IPv4`.

Root Cause a self-feedback loop, NOT PF. `scripts/watcher.sh` Listener 2 watches `State:/Network/Global/IPv4` via `scutil n.watch` and launches `recover.sh --post-wake` on any change. But post-wake re-asserts the exit node (`recover.sh` ~L367, `tailscale set --exit-node`), and even though `set` is deliberately used instead of `--reset` to minimise churn (recover.sh ~L360-366), **re-asserting the exit node re-writes the host default route and the default route IS `State:/Network/Global/IPv4`.** So each run mutates the exact key the watcher triggers on. The self-triggered change lands ~15-45s after the run finishes; `run_recover()`'s old 10s debounce (reset to the run's *finish* time at the tail of the function) had long expired by then, so the watcher re-fired. The loop self-terminates only once `tailscale set` becomes a no-op (route already correct) no route change no re-trigger.

Why the debounce couldn't stop it: the loop period (~40s = one full post-wake run + settle) far exceeds the 10s debounce, which therefore only ever caught macOS's rapid *duplicate* scutil notifications for a single change (the `+0.2s` doublets), not the self-trigger.

Not PF: post-wake never calls `pfctl` (the kill-switch is armed once, in `scripts/common.sh`, not per cycle); no PF-enable line falls inside the loop window. The `No ALTQ support in kernel` warning is normal macOS pfctl noise (Apple ships pfctl without ALTQ compiled in), and `Could not capture PF reference token` is a separate benign ref-count bookkeeping wart both unrelated to the loop.

Fix: `scripts/watcher.sh` `run_recover()` now arms a **post-recovery settle cooldown** (`/tmp/nullexit-watcher.settle-until`, `NULLEXIT_POSTWAKE_SETTLE_SECONDS` default 45s) *after* each post-wake run, and skips launching while `now < settle-until`. Because it is armed only after a recovery runs, a genuine first wake/roam is still handled immediately; only the Global/IPv4 change the run caused itself is absorbed. A real new roam during the window is delayed at most `POSTWAKE_SETTLE_SECONDS`, an acceptable trade for killing the loop. Watcher-side only `recover.sh` is unchanged. Severity was low (self-correcting, healthy every iteration, no leak/outage; cost was wasted ~28s runs + log noise).

15.12.22 FaceTime (and Real-Time P2P Media) Over the Double-Tunnel: Why It Breaks + the Split-Route Plan (July 15, 2026)

Symptom: A FaceTime call placed while nullexit is up fails to connect.

Diagnosis nullexit is NOT blocking it; every layer passes the traffic (all verified live):

- **DNS:** no Apple/FaceTime domain is sinkholed `dig` via the gateway (`100.100.21.8`) returns real Apple A/CNAME records, none `0.0.0.0`/`127.0.0.1`. The `<no answer>` domains (`gateway.push.apple.com`, `fmn.apple.com`, `identityservices.apple.com`, `fdr.apple.com`) return empty on `1.1.1.1` too CNAME/non-A, not our block.
- **routing-fix `FORWARD`:** the exit-path rule `-i tailscale0 -o tun0 -j ACCEPT` carries ~203K pkts / 73M bytes and accepts *everything*.
- **Country / threat blocklists** (`block_il`, `block_kp`, `block_malicious`): **0** packets dropped.
- **WARP (gluetun) egress:** `-A OUTPUT -o tun0 -j ACCEPT` + `-A POSTROUTING -o tun0 -j MASQUERADE` all outbound UDP (any port) leaves and is SNAT'd through WARP.

Root cause double symmetric NAT ending at Cloudflare WARP. Path: `host Tailscale (MASQUERADE) WARP (MASQUERADE) Cloudflare shared egress`. WARP is a CGNAT-style shared egress with **no inbound UDP mappings**, and the double-MASQUERADE presents as **symmetric NAT** the exact condition that defeats FaceTime's STUN/ICE UDP hole-punching. No P2P path can form; the Apple-relay fallback is unreliable through the stacked tunnel, and the 12801200 MTU squeeze degrades media. This is inherent to routing real-time P2P media through WARP, **not** a firewall rule. (Zoom/WhatsApp *calls* hit the same wall.)

Latent finding (separate bug) Closed: the `FORWARD` chain's `-i tailscale0 -p udp --dport 443` (QUIC) and `--dport 853` (DoT) `REJECT` rules were **shadowed** by the earlier `-i tailscale0 -o tun0 -j ACCEPT` (first-match wins) verified **0 pkts** on each at the time. Fixed by moving those port-`REJECT` rules in `post-rules.txt` **before** the `tailscale0 tun0` / `tun0 tailscale0` ACCEPT rules. Re-verified live post-fix: the REJECTs now sit ahead of the ACCEPTs in `iptables -L FORWARD --line-numbers` and accumulate hits.

Fix plan (chosen: Option 2 narrow, always-on, opt-in split-route). Implemented and shipped enabled by default:

- **Mechanism:** add more-specific host routes for FaceTime's Apple `/16`s via the **physical** gateway (`172.17.0.1`/`en0`). Longest-prefix match beats Tailscale's exit-node `0.0.0.0/1` + `128.0.0.0/1` routes, so only those `/16`s leave **direct** while everything else stays double-tunneled. Mirrors `add_warp_bypass_routes`.
- **Wiring:** `add_apple_split_routes()` / `remove_apple_split_routes()` in `common.sh`; called in toggle START (after exit-node assert) + `recover.sh --post-wake` (survives roam) + removed on STOP; guarded by `FACETIME_SPLIT_ROUTE` in `.env` `.env.example` ships it `false` (opt-in, privacy-first default for new installs), enabled here after the tradeoff below was accepted and the CIDR verified against a real call.
- **Privacy tradeoff:** whatever `/16`s go direct see the **real ISP IP** (not WARP). Apple already ties traffic to the Apple ID, so Apple-direct is the accepted scope; user chose **narrow** (FaceTime-infra `/16`s only) over all-`17.0.0.0/8` to minimise the surface.
- **CIDR list finalised from an empirical capture:** host-side `sudo tcpdump -n -i <exit-utun, e.g. utun6> net 17.0.0.0/8` during a real FaceTime call. The capture identified **`17.249.0.0/16`** as the load-bearing media/relay range this is the value shipped in `.env.example`'s `FACETIME_DIRECT_SUBNETS`, verified working end-to-end on a real call. (Earlier DNS-only guesses `17.57/16` APNs push/ringing, `17.248/16` iCloud/FaceTime gateway, `17.253/16` aaplimg push CDN, `17.157/16` GSA auth were signaling-only ranges and did not carry media; the capture superseded them.) See P2P note 16.4.

15.12.23 setup.sh Never Generated NULLEXIT_SEED Every Fresh Install Hard-Failed Its First toggle.sh (July 18, 2026)

Symptom: README "Cryptographic Integrity Verification" and this doc's own `.env` variable table both claimed `NULLEXIT_SEED` is "generated by `setup.sh`." It wasn't. `scripts/setup-common.sh`'s `.env`-writing step never referenced `NULLEXIT_SEED` at all, and `.env.example` only offered a placeholder telling the user to run `openssl rand -hex 32` by hand. `.signatures` is committed to git (tracking the maintainer's own local seed), so a fresh clone's `toggle.sh` would call `scripts/crypto.sh --verify`, find no `NULLEXIT_SEED` in the new `.env` at all, and hard-fail with a cryptic seed error on the very first run with nothing in the setup flow or its closing instructions telling the user a `crypto.sh --sign` step was needed, or why.

Compounding bug found in the same code path: `setup-common.sh`'s `.env` writer also defaulted `KILL_SWITCH` to `"false"` for brand-new installs (`EXISTING_KILL="false"`, only overridden if an `.env` already existed with a different value) the exact inverse of the "leave true, this is the core leak guarantee" invariant enforced everywhere else (`.env.example`, `scripts/profiles.sh`). Every fresh install would have shipped with the kill-switch off by default, not just failed on first boot.

Fix: `setup-common.sh` now generates `NULLEXIT_SEED` via `openssl rand -hex 32` when `.env` doesn't already have one (preserving it on re-runs, same pattern as every other preserved flag), writes it into the new `.env`, and calls `scripts/crypto.sh --sign` immediately after writing `.env` so a fresh install's first `toggle.sh` finds a valid seed and matching signatures instead of failing cold. `EXISTING_KILL` now defaults to `"true"`.

15.12.24 Honey-Port Subshell Holds stdout/stderr Open Piping `toggle.sh` Hangs Long After It Finishes (July 18, 2026)
Symptom: `./toggle.sh --restart 2>&1 | tail -120` appeared to hang indefinitely 10+ minutes with zero output even though `output.log` showed the restart itself completed cleanly (`EXIT: code=0 phase=start`) in about a minute and all 5 containers were healthy. `ps` showed the top-level `bash ./toggle.sh` process still alive, reparented to launchd (`ppid=1`), sitting in a `__wait4` syscall (confirmed via `sample`) with its only child being the [[Honey-Port Tripwire]]'s `nc -l localhost <port>` listener.

Root cause: the tripwire's arming code (`toggle.sh` ~L1704) redirects only the *inner* `nc -l` command's own stdout/stderr to `/dev/null` the **subshell process itself** (`( if nc -l  ; then  fi ) &`) inherits fd 1/2 unchanged from its parent. Since `nc -l` never returns (by design it's waiting for a port-scan that almost never comes), that subshell stays alive indefinitely, and because it still holds the *write end* of whatever pipe its parent's stdout was connected to, any reader on the other end (`| tail`, `| tee`, a CI log capture, anything) never sees EOF even though the disowned process is functionally idle and the script that spawned it exited normally minutes earlier. This is a distinct manifestation of the listener-lifetime issue already known from 15.12.20 (accumulation): that fix (`pkill -f "nc -l localhost"` before arming) reaps the *previous* listener, but does nothing about the *current* one's inherited file descriptors.

Workaround used: killed only the reader (`tail`) on the stuck end of the pipe safe, since it doesn't touch the gateway, Docker, `caffeinate`, or the tripwire itself, which is meant to keep running. **Not yet fixed at the source** see [[Future - Pluggable Egress|18 TODO]] candidate: redirect the subshell's own stdout/stderr (`(  ) >/dev/null 2>&1 &`) at arm time so it can no longer hold a caller's pipe open. Already flagged as a sharp edge in the [[toggle.sh]] component doc ("Never pipe it through a reader") this incident is the first confirmed real-world hit, not just a theoretical gotcha.

15.12.25 Incident Post-Mortem: WARP Watcher False-Kills the Gateway on a Wi-Fi Drop; Wi-Fi Auto-Rejoin (July 2223, 2026)
Symptom: Overnight the gateway went fully down and stayed down (~22:46 EDT until it was manually brought back). No mesh device including the S24 had connectivity during the window. The user flagged it as suspicious ("what killed the watcher").

Root cause (not malicious): the in-container WARP Watcher (15.6.2) polls `cdn-cgi/trace` and cannot distinguish "WARP tunnel failed" from "the host physical link (en0) is down" both surface as `warp=off`, because with no internet there is no tunnel *and* no trace response. The host lost Wi-Fi (`ipconfig_get_packet(en0) failed: not found` repeated in `output.log`), the watcher counted its 6 consecutive `off` polls, and `recover.sh` tore the whole gateway down a gateway that would have healed on its own once Wi-Fi returned. A secondary design gap amplified it: `recover.sh --post-wake` releases the watcher inhibit marker on a *timeout* (after `wait_for_dhcp_settle` gives up) rather than waiting for `warp=on`, so on a drop the inhibit lifted while WARP was still legitimately down.

Fix, part 1 pause, don't kill. The WARP Watcher now checks the physical uplink (`ipconfig getifaddr $phys_iface`) before treating a `warp=off` as a tunnel failure; empty/`169.254.*` PAUSE (reset counter, sleep, `continue`) instead of running `recover.sh`. Leak-safe by construction (no routable IP = no egress path). Full detail in 15.6.2 point 5.

Fix, part 2 actually reconnect the Wi-Fi. Pausing stops the false kill but the Mac still has no internet if macOS auto-join is off; it "just sits there." macOS (Darwin 25) redacts the live SSID from every script unless the process holds Location, and the user explicitly refused Location (15.11.10), so the *current* network is undetectable. Solution: don't detect it **remember** it. `scripts/wifi-rejoin.sh` keeps a user-managed list of network names in `.last_wifi_ssid` (`set`/`forget`/`list`, gitignored, unsigned) and re-associates the first remembered one via `sudo -n networksetup -setairportnetwork "$iface" "$ssid"`, which pulls the WPA password from the login Keychain automatically (confirmed by live drop test) and only ever targets a network the user explicitly remembered no scanning, no public-AP hopping. `rejoin` arms a down-timer on first call, holds a 15s **grace** period (let DHCP self-heal a same-AP blip first), then attempts with **progressive backoff** (3060120240cap 300s) and **gives up after 6 tries** per down-episode so it never churns a network that can't come up (captive portal, DHCP blocked, out of range); `reset` clears the timer once the link returns.

Fix, part 3 put the driver where it can't die with the gateway. The first cut wired `rejoin` into the WARP Watcher only but that watcher runs *solely while the gateway is up*, so a Wi-Fi drop while the gateway was down/off left nothing to reconnect the Mac (observed directly: a retest "just didn't do shit" because the gateway had shut down at 06:54:55Z and its watcher was gone). The rejoin driver was therefore added as a 4th listener in the **always-on** `scripts/watcher.sh` daemon (15.5.9): a 15s poll of the Wi-Fi interface that calls `wifi-rejoin.sh rejoin` when it has no routable IP and `reset` when it does. The WARP Watcher still calls the same helper while the gateway is up; the two coordinate through shared `/tmp` markers rather than double-associating. Net result: Wi-Fi reconnects whether the gateway is up, down, or off which is what the user's original report ("sometimes Wi-Fi disconnects and it just sits there without internet") actually required.

Self-inflicted detour worth recording: an intermediate retest broke the network *after* reconnecting because editing signed core files (`toggle.sh`, `docker-compose.yml`) without re-running `./scripts/crypto.sh --sign` made `recover.sh --post-wake` fail its integrity check and abort (exit 1) so recovery never ran and `wifi-rejoin` churned every 30s. Lesson reinforced: **re-sign after every edit to a signed file** ("Cryptographic Integrity Verification"), and the give-up cap above exists partly to make such a churn self-limiting.

15.12.26 Logging Unified to UTC (July 23, 2026)
Symptom: `output.log` interleaved UTC and local (EDT) timestamps, and the AdGuard query log was in local time while `blocked.log` was UTC making forensics (e.g. "was the S24 connected at 23:00?") error-prone, because the same wall-clock instant appeared under two different hour labels. The user asked for "one thing, utc or est, only one."

Root cause: `TZ=America/New_York` was set on every service in `docker-compose.yml`, but it is only honored by images that ship `tzdata` (AdGuard rendered EDT); Alpine-based images without `tzdata` (routing-fix, the Go logger) silently fell back to UTC. So the timezone of a line depended on which container emitted it.

Fix: standardized on **UTC everywhere**. All `TZ=` entries in `docker-compose.yml` were set to `UTC`, and `scripts/logger/main.go` now formats with `time.Now().UTC()` and appends an explicit ` UTC` marker to each `blocked.log` line so the zone is unambiguous on the page. `docker-compose.yml` is a signed file and was re-signed. Host scripts (`watcher.sh`, `wifi-rejoin.sh`, `recover.sh`) already emit UTC via `date -u`. **Forensics caveat:** log lines written *before* 2026-07-23 still carry the old UTC/local split see the [[log-timezones]] note when reading historical entries.

15.12.27 Incident Post-Mortem: WARP Watcher False-Kill Recurred "Has an IP" "Has Egress" (July 23, 2026)
Symptom: the gateway tore itself down at **12:0312:04 EDT (16:0316:04Z)** and stayed dead for **2 h 41 m** until a manual `toggle.sh` at 14:45 EDT even though the 15.6.2 point-5 pause guard, shipped that very morning, was supposed to prevent exactly this. User report: *"a few hours ago nullexit stopped without any reason."*

Timeline (from `output.log`):
```
16:03:08Z  WARP DOWN  cdn-cgi/trace reports warp=off       (en0 still shows a leased IP)
16:03:59Z  WARP SHUTDOWN  6 consecutive failures  recover.sh --auto  (whole gateway nuked)
16:04:04Z  recover.sh power-cycles Wi-Fi  security='none' (no-wifi)
16:04:21Z16:22:01Z  wifi-rejoin attempts 1..6 all fail  "gave up, reconnect manually"
18:45:02Z  (14:45 EDT) user manually re-runs toggle.sh; gateway restored
```

Root cause: the guard's uplink test was `ipconfig getifaddr $phys_iface` alone, and *having an IP is not the same as having egress.* The host was on **`ZUM-WiFi`, an 802.1x/WPA2-Enterprise network** (`P2P detect:  security='Enterprise'` throughout). On an enterprise de-auth/roam, macOS keeps the **stale DHCP lease** on `en0` for a while after the association is already dead so `getifaddr` returned a real, non-`169.254` address and the guard concluded "uplink present," skipped the pause, and let the failure counter run. The container's `cdn-cgi/trace` probe *correctly* reported `warp=off` (there was genuinely no internet), 6 polls accumulated in ~50 s, and the nuclear `recover.sh` fired. This was **not** a real WARP-tunnel failure: there was no "warp container tunnel is dead" force-recreate anywhere near 16:03 (the only one that day was the unrelated 02:51 EDT event), and the fact that all **6 `wifi-rejoin` attempts failed over 18 minutes** proves the *network itself* was gone (a live AP re-associates on attempt 1) i.e. the trigger was an uplink loss misread as a tunnel death, the same class as 15.12.25. Enterprise Wi-Fi made both halves worse: it widens the stale-lease window, and with macOS auto-join off the 802.1x re-auth never completed on its own, so the gateway couldn't self-heal.

Fix: add a **link-state check** to the guard (full detail in 15.6.2 point 5). After confirming a routable IP, the watcher now reads the driver's association state via `ifconfig $phys_iface` `status: active` proceeds to the WARP check, `status: inactive` (stale lease over a dead link) PAUSE. A ping-to-router probe was tried first and **rejected after live testing**: the PF kill-switch + this network's AP isolation drop ICMP to the LAN gateway even when `warp=on`, so a ping guard would false-PAUSE a *healthy* gateway worse than the original bug. The link-state read is ICMP/PF-independent and fails open (unknown status trust the lease). PF kill-switch remains the leak guarantee, so a wrong pause is "no internet," never a real-IP leak. `toggle.sh` re-signed (`scripts/crypto.sh --sign`); takes effect on the next `./toggle.sh --restart`.

Forensics caveat: reconstructing this was slowed by a residual timezone split the 15.12.26 unification did **not** cover the container/`date -u` lines are UTC (`...T16:03:08Z`), but `toggle.sh`/`recover.sh`'s coloured `step()`/`ok()` lines still print **local EDT** in the `[2026-07-23 12:03:59]` bracket-space format (the same POST-WAKE instant appears as both `12:03` and `16:03Z`). 15.12.26 fixed the container logs and the Go logger but not the bash echo helpers; when reading a single incident, convert bracket-space stamps as EDT = UTC4. See [[log-timezones]].

Part IV Observations, Changelog & TODO

16. Unverified Observations

16.1 AdGuard Returns Two Different Sinkhole IPs (Unverified)

> **Status: Observed but not formally verified. Needs a deeper audit of the rule-compiler sources to confirm.**

Observation (July 10, 2026): When querying AdGuard Home for blocked ad domains, some domains resolve to `0.0.0.0` and others resolve to `127.0.0.1`. Both are blocked, both cause connection failure, but the different responses were unexpected.

```
doubleclick.net              0.0.0.0    (blocked)
ads.google.com               127.0.0.1  (blocked)
pagead2.googlesyndication.com  0.0.0.0  (blocked)
scorecardresearch.com        127.0.0.1  (blocked)
```

Likely Explanation (Unverified): There are three historical schools of thought on the "correct" DNS sinkhole response, and AdGuard faithfully preserves whichever the source blocklist used:

```
| Sinkhole IP | Convention | Origin |
|---|---|---|
| `0.0.0.0` | AdBlock-style lists (`\|\|domain^`) | Modern DNS-level blocking; AdGuard's native format. Fails fast  OS doesn't attempt a TCP connection. Works for both IPv4 and IPv6 without a separate rule. |
| `127.0.0.1` | Hosts-file-style lists (`127.0.0.1 domain`) | 90s-era `/etc/hosts` tradition. Steven Black's hosts list (500k+ domains) still uses this. |
| `NXDOMAIN` | Purist / Pi-hole default | Returns "domain doesn't exist." Semantically most accurate but requires browsers to handle NXDOMAIN gracefully. |

The dual-response behavior in this project is because `rule-compiler` pulls from both AdBlock-format sources (Hagezi, uBlock origin lists) and hosts-format sources (Spamhaus DROP, Steven Black, Feodo Tracker), and AdGuard returns whatever sinkhole IP the matching rule specifies.

**To Verify:** Run `docker compose exec tailscale cat /etc/hosts` to confirm no hosts-file rules are being injected at the container level, and check which specific AdGuard rule matched each domain via the AdGuard query log at `http://100.100.21.8:3000/#logs`.

### 16.2 AGY CLI Appears to Generate Native macOS Notification Pop-ups (Unverified)

> **Status: Observed but source unverified. macOS security logs ruled out system-level origin.**

**Observation (July 10, 2026):** While running a DNS blocklist verification via the Antigravity CLI (`agy`), a command was proposed that included `malware.wicar.org` (a known-safe DNS blocklist test domain) in a `dig` loop. Shortly after, a macOS-style notification popup appeared describing a "malicious" or "blocked" event. The AGY CLI terminal session then closed.

**Investigation:**
- **macOS XProtect / Gatekeeper logs:** Empty. Zero events in the 30-minute window around the incident.
- **macOS Endpoint Security / System Policy logs:** Empty.
- **Broad `log show` search for "malicious", "malware", "blocked":** Empty.
- **Conclusion:** macOS itself did not generate the notification. No system security subsystem fired.

**Likely Explanation (Unverified):** The notification appears to have originated from **AGY CLI's own internal safety system**. AGY CLI is a native macOS application (not just a terminal process) and has access to the macOS Notification Center API (`UNUserNotificationCenter`). When its guardrails detected a domain containing the word "malware" (`malware.wicar.org`) in a proposed shell command, it likely:
1. Blocked the command from executing
2. Posted a native macOS-style notification via the Notification Center

This is surprising behavior for a CLI tool  most terminal programs don't send GUI notifications  but AGY CLI appears to bridge both worlds: it runs in a terminal but is packaged as a native app with full system API access. The notification would be indistinguishable from a system security alert at a glance.

**Notes:**
- `malware.wicar.org` is the DNS/AV equivalent of the EICAR test file  a deliberately benign domain that triggers blocklist/AV systems to verify they work. It was included in the test set to confirm the AdGuard blocklist was catching it. AGY's safety system is not aware of this distinction.
- For future DNS blocklist testing, use only ad/tracking domains (e.g., `doubleclick.net`, `adnxs.com`) to avoid triggering AGY's guardrails.

### 16.3 `system_profiler SPAirPortDataType` Longevity (Wi-Fi Detection)
The Wi-Fi security detection used by `watcher.sh`'s `detect_lan_p2p_mode()` currently relies on `system_profiler SPAirPortDataType` (after the `airport` binary removal  full history in 15.11.6). This works as of macOS Sequoia (15.x) without sudo and without redaction, but given Apple's ongoing Wi-Fi privacy clampdown, it may be locked down in a future release. If `security` comes back empty on a future macOS version while on Wi-Fi, detection silently falls back to `allow=false` (a safe default). The long-term fix would be a small Swift CoreWLAN helper compiled and bundled in the repo.

### 16.4 Tailscale LAN P2P Succeeds Cross-AP But Fails Same-AP (Client Isolation + No NAT Hairpin)

**Observation (July 15, 2026):** On the residence Wi-Fi, a Galaxy S24 in the lobby reception  6 floors away, associated to a *different* AP  establishes a **direct, low-ping** Tailscale P2P session to an iPhone 11. The *same* pair on the *same* AP / same floor / close range does **not** go direct. Counterintuitively, **farther apart = P2P works**. The effect shows for the iPhone 11 (a plain client) but **not** for S24macOS or S24the gateway container.

**Mechanism:**
- **Same AP:** the AP enforces **client isolation** (L2 devicedevice block on one radio). Both devices then share one public IP; STUN returns identical external mappings, so reaching each other would need the residence router to **NAT-hairpin**  which most campus/consumer gear does not do  direct P2P fails and falls back to a DERP relay.
- **Different AP (floors away):** large residence networks **segment per AP/area into distinct VLANs/subnets**, so the two devices sit on different L3 networks routed through the campus core, where client isolation does not apply and the STUN path works **without** hairpin  clean direct P2P.
- **Why macOS / the container are immune:** nullexit pins a **direct WireGuard endpoint** (Tailscale port `41642`, `direct 172.x`), so those peers already hold a negotiated direct path and do not depend on same-AP L2 discovery. A plain iPhone 11 has no pinned endpoint, so its path is entirely at the mercy of the Wi-Fi topology  which is exactly why *it* exposes the isolation effect.
```

Relevance: refines `watcher.sh`'s `detect_lan_p2p_mode()` (which probes same-subnet AP isolation via a broadcast ping). The probe correctly flags same-AP isolation, but this observation adds that isolation is **per-AP** and cross-VLAN P2P can succeed even when same-AP fails so a single "LAN P2P allowed?" boolean under-describes a multi-AP network. No code change; documents expected behavior. Related: the FaceTime NAT analysis in 15.12.22 (same STUN/hairpin physics, different symptom).

Follow-up (July 18, 2026) measuring the same-AP DERP jitter, and whether P2P can be forced anyway:

Burst-pinged (`tailscale ping --c 8`) an S24 on this exact same-AP-isolated network: 8/8 pong, 0% loss, but RTT swung 38-703ms (665ms spread) over `DERP(tor)`. The gateway's own path to the same Toronto DERP relay measured a clean, stable 29.1ms (`tailscale netcheck`), ruling out anything on the nullexit/gateway side the variance is entirely downstream, between the S24 and the relay.

Switched the S24 to cellular (same physical location, different link) for a controlled comparison: 78-500ms (422ms spread) somewhat better, but still substantial. If the jitter were purely Android's Wi-Fi power-saving radio behavior, cellular should have looked dramatically cleaner; it didn't, which points at the real shared cause instead: both Wi-Fi and cellular radios drop to a low-power idle state between packets and pay a wake-up latency penalty on each new one, and DERP relaying rides over TCP (its own retransmit/ACK timing adds further variance that's invisible as "loss" at the WireGuard layer). This is inherent to relaying sparse/bursty traffic over *any* mobile radio link not a nullexit configuration issue, and not fixable from the gateway side.

Can P2P be forced anyway, despite confirmed AP isolation? Confirmed with the user this specific network genuinely has AP isolation (not a `detect_lan_p2p_mode()` false positive). Per the mechanism above: the block is enforced by the AP at Layer 2, below anything WireGuard/Tailscale operates at no VPN-layer trick un-isolates two clients the radio itself refuses to bridge. The only theoretical bypass is the residence router performing **NAT hairpin** (loopback) so both devices' identical STUN-mapped public IP routes back down through the router instead of laterally through the AP and this section already established "most campus/consumer gear does not do" that. AP-isolated deployments are, if anything, *less* likely to support hairpinning than typical consumer gear, since both are usually part of the same client-isolation security posture. **Conclusion: no, not reliably forceable.** `TAILSCALE_ALLOW_LAN_P2P=true` would not restore P2P here it would only reproduce the 15.7 SNAT endpoint-poisoning blackhole, trading "jittery over DERP" for "briefly offline entirely." Leave it `false`. (One untested, unguaranteed exception: some AP isolation implementations only block wirelesswireless traffic at the radio level, not wiredwireless a **wired** Ethernet connection for one side might bypass it on hardware that isolates that way. Depends entirely on how this specific AP implements isolation; not something to assume works.)

16.5 Working Theory iPhone Cross-Device Visibility Reaches Other Phones, Not macOS (Instagram Auto-Open, July 15, 2026)

Observation: Opening Instagram on a Galaxy S24 (specifically via **Chrome** on the S24, not Samsung Internet) auto-opens a tab on a nearby MacBook, but **only when the S24 is physically close** to the Mac it does not happen when the phone is away. **Chrome on the S24 has zero granted Android permissions** (no Bluetooth, no location, no nearby-devices) this rules out Chrome itself performing any BLE/Nearby proximity scan, since Android enforces those permissions at the app level regardless of what the browser ships. This pushes the likely actor down to **Google Play Services**, which holds system-level Bluetooth/nearby-device access independent of per-app grants and could be doing the proximity detection on Chrome's behalf (e.g. via Android's "Nearby" APIs) then handing the result to Chrome only for the actual tab-open. A purely local, short-range channel is still required to explain the proximity gating regardless of which component performs it; this also rules out any cloud-mediated mechanism (Chrome cloud tab sync, FCM push, Tailscale/WAN P2P per 16.4) since those are distance-independent. Live probes during the session (`system_profiler SPBluetoothDataType` for a paired device, `dns-sd -B _services._dns-sd._udp` for 8s) found **no paired Bluetooth device and no Android/Google/Instagram mDNS service** only the Mac's own Bonjour services (AirPlay, RFB, `_companion-link`, Spotify Connect). Inconclusive: a BLE-advertisement-only channel (no formal pairing) would not show up in either probe.

User's working theory: the iPhone (also on hand, on the same network/Apple ID as the Mac) makes *other devices connected to it* even an Android phone visible to nearby devices, but this visibility channel reaches **other phones more readily than it reaches macOS**. I.e., the iPhone may be acting as a discovery bridge/relay between the S24 and the Mac rather than the S24 and Mac talking directly. Combined with the Play Services angle above, a refined version is: **Play Services on the S24 detects proximity to an Apple mesh (iPhone and/or Mac both signed into iCloud/AWDL-visible) via some cross-ecosystem discovery Google has built**, rather than the S24 and Mac negotiating directly.

Cross-ecosystem refinement: this isn't Google-only. Apple runs an equivalent OS-level background BLE scanning layer for its own Continuity stack (Handoff, Universal Clipboard, AirDrop discovery, Find My the latter using *every nearby Apple device*, not just the owner's, as a passive BLE relay), implemented as system daemons (`bluetoothd`/`sharingd`/`rapportd`, already implicated in the unrelated AWDL wedge bug at 15.10) rather than anything gated by app-level permission. The key mechanism that makes the two ecosystems mutually visible: **BLE advertisement packets are broadcast in the clear over open air** Apple's Continuity BLE frames are semi-proprietary but unencrypted at the advertisement layer, so any BLE radio (including Android's, via Play Services once Android's own "Wi-Fi & Bluetooth scanning" system-level toggle is on) can passively observe "an Apple device is here" and react, without pairing or explicit permission from Apple's side. So both ecosystems maintain a BLE mesh-discovery layer that sits below their own per-app permission model, and each side's layer is passively readable by the other's radio which is a plausible bridge for a Play-Services-mediated Chrome trigger reacting to iPhone/Mac Continuity broadcasts.

Status: Unverified. Not yet reproduced under an active capture the July 15 session ran the mDNS browse *before* confirming the phone was actively broadcasting, so it's not a clean negative result. To confirm: capture `tcpdump -i en0 udp port 5353` (mDNS) **and a raw BLE advertisement capture** (e.g. `nRF Connect` or similar BLE scanner app on a third device, or `sudo btmon`-equivalent on Linux macOS does not expose

raw BLE advertisement capture without a hardware sniffer) at the exact moment Instagram is opened on the S24 while close, with the iPhone present vs. absent, to isolate whether the iPhone's Continuity broadcasts are the trigger or the S24Mac channel exists independently of it. On the Android side, checking `Settings  Location  Location Services  Wi-Fi & Bluetooth scanning` and whether **Google Play Services** (not Chrome) holds Bluetooth/nearby-device/location permissions would confirm or kill the Play-Services-as-actor theory. Related: 16.4 (same physical-proximity family of P2P/discovery quirks, different mechanism and inverted distance relationship).

17. Changelog / Recent Updates

Reflects the current branch's state. Each entry is a one-line summary + commit hash; read the commit message (and the linked Incident Log entry) for full detail. Full post-mortems live in 15.

- **July 23, 2026 `fix(warp-watcher): link-state check closes the stale-lease false-kill`** The 2026-07-23 pause guard only checked `getifaddr`, so an 802.1x de-auth that left `en0` holding a stale DHCP lease still let the watcher nuke the whole gateway (2 h 41 m outage on `ZUM-WiFi`). The guard now also reads the NIC's `ifconfig` association state (`status: active`) before counting `warp=off` as a tunnel failure "has an IP " is not "has egress." (A router-ping probe was tried and rejected: PF + AP-isolation drop ICMP even when healthy, so it would false-pause a live gateway.) Fails open; PF kill-switch stays the leak guarantee. See 15.12.27 and 15.6.2 pt 5.
- **July 19, 2026 docs: dense-cellular field test** Live burst-ping test of a mesh device on cellular in a very dense crowd: 0% loss, DERP-relayed the whole time (expected on cellular NAT, not congestion-related), a couple of RF-driven latency spikes traced to the cell tower rather than the gateway. See 15.3.7.
- **July 18, 2026 `fix(toggle): harden sleep prevention with caffeinate -s`** Four real `Idle Sleep` cycles in a 38-minute window blacked out every mesh device's exit node because `caffeinate -i` alone doesn't reliably hold sleep off. Added `-s` (blocks sleep on AC power, no-op on battery) alongside the existing `-i `. See 15.5.8.
- **July 18, 2026 docs: honey-port pipe-hang incident** Confirmed in practice (not just the known theoretical gotcha) that the Honey-Port subshell's inherited stdout/stderr keeps any `| reader` on `toggle.sh` blocked long after the script itself exits. Worked around by killing the reader; source fix still open. See 15.12.24.
- **July 10, 2026 `fix(race): suppress WARP Watcher during post-wake force-recreate`** Fixed a race where switching Wi-Fi caused the host IP to leak as the raw ISP IP (`132.x`) on every roam; `recover.sh --post-wake`'s warp force-recreate tripped the WARP Watcher into a nuclear shutdown. Fixed via `/tmp/nullexit-warp -inhibit.marker`. See 15.2.7.
- **July 10, 2026 startup/roaming stabilization suite** Pre-flight Tailscale race (15.5.3), STUN-block cold-boot false negative (15.5.4), restart DNS-build race via `wait_for_dhcp_settle` (15.5.5), nuclear/post-wake concurrency lock (15.2.8), infinite auto-recovery loop marker leak (15.2.9), PF kill-switch bricking SOCKS5 fallback (15.7.3), `--reset` host logouts (15.7.4), missing-flags abort (15.7.5), discarded docker-exec errors (15.12.18), and the SNAT endpoint poisoning fix (15.7.1).
- **July 6, 2026 `fix: resolve edge-case vulnerabilities and system state leaks`** Security-review hardening suite:
  1. **DNS Watcher Interface Independence:** `start_dns_watcher` now detects the active interface via ` get_active_service()` instead of hardcoding a `grep` for "Wi-Fi".
  2. **Robust Lock Verification:** `toggle.sh` lock-file logic uses `ps -p` verification to prevent permanent lockouts from a recycled PID.
  3. **Fail-Closed Crypto:** `crypto.sh` integrity check switched from `[ -x ]` to `[ -f ]` so it runs even if git strips the `+x` bit.
  4. **Strict `iptables` Pinning:** `post-rules.txt` FORWARD-chain rules explicitly pin `tailscale0``tun0` both  directions, preventing escape via `eth0` during WARP drops.
  5. **Docker Local Network Isolation:** `docker-compose.yml` binds exposed WARP ports (`1080`, `5354`) to `127.0.0.1` to prevent LAN access.
  6. **Routing Verification & Sudoers:** Added a `netstat -nr` default-route override check and ensured `/usr/bin/python3` is permitted in the sudoers template.
- **July 4, 2026 `cc789c5`** Concurrency mutex on `toggle.sh` (`/tmp/nullexit-toggle.lock`); defensive ` TS_AUTHKEY` init under `set -u`; Go `go vet`/`go build` steps in CI; AdGuard log capping; macOS BSD `grep` fixes in `setup-common.sh`; `recover-linux.sh` quoting/PID fixes; removed redundant bypass-routes block in `recover.sh`.
- **July 4, 2026 `fix: post-wake routing loop & destructive recovery`** (1) `add_warp_bypass_routes` now accepts explicit interface args so post-wake stops routing WARP traffic back into `utun*` (was an infinite loop killing internet on every wake). (2) Replaced `sudo route -n flush` with targeted deletions of `0.0.0.0/1`/`128.0.0.0/1` + WARP bypass routes, so recovery no longer destroys loopback/multicast routes and wedges `configd`/`mDNSResponder` until reboot.
- **July 4, 2026 `feat: secure execution`** Restricted the blanket `/usr/bin/python3` NOPASSWD sudo to exactly `/usr/bin/python3 -I .../scripts/dns-proxy.py`; locked `dns-proxy.py` with `chown root:wheel` + ` chmod 755`; added native auth prompts to `Toggle Gateway.app`/`Recover Gateway.app`; `toggle.sh` now captures `warp` logs on failure before teardown.
- **July 4, 2026 Go ports + refactor** `logger` and `rule-compiler` ported to Go multi-stage Docker builds (15.8.6); ~200 lines of duplicated bash extracted to `scripts/common.sh` (15.9.4); `crypto.sh` HMAC-SHA256 script integrity added.
- **July 3, 2026 `fix: resolve numerous system regressions`** Fixed truncated `/etc/sudoers.d/nullexit` line; temporary `.tmp`+`os.replace()` atomic writes in `sync-rules.py` (later reverted, see below); `recover.sh` `ts_args` quoting; ANSI `%b`/`-e` output in `diagnose-host-leak.sh`/`fix-docker-bridge-collision.sh`; ` toggle-linux.sh` using `resolvectl dns` instead of macOS `networksetup`; AdGuard REST refresh POST targeting port `3000`; restored `START_GATEWAY=true`; purged stale port `80` mapping; verified `output.log` in `.gitignore`.
- **July 3, 2026 `unlock-files.sh` + revert atomic writes** Reverted all 5 `.tmp`+`os.replace()` sites in ` sync-rules.py` to direct writes; created `scripts/unlock-files.sh` to permanently fix stale `0444`/`000` file permissions via atomic rename (no `chmod`). Covers `.env`, `compiled_rules.txt`, `ip_blocklist.ipset`, `cache/*.txt`, `cache/ip/*.txt`, `data/filters/*.txt`. See 15.9.6 + 3.

2229 - **July 3, 2026 Overnight leak + geo-blocking** Introduced `scripts/host-leak-probe.sh` (sub-second host-egress detector, 15.6.1) and the statistical-threshold WARP auto-shutdown watcher (15.6.2); fixed the overnight silent IP leak (15.6.3).
2230 - **July 2, 2026 `8c5fae1` + `181e1e1`** Two infinite routing loops fixed: host-VM tunnel loop (WARP bypass routes via physical gateway IP + `setup_exit_node_routing` forcing default to `utun*`) and Docker Compose subnet takeover (`10.200.1.0/24` lock). `--accept-routes=true` added explicitly to all `tailscale up --reset` calls. See 15.2.2.
2231 - **July 2, 2026 auto-recovery daemon** `scripts/watcher.sh` + launchd LaunchAgent documented. See 15.5.1.
2232 - **July 1, 2026 `c5b92c0` (refactor: personal-leak cleanup)** Eliminated every residual personal-username leak across source + config; launchd plist uses `WorkingDirectory` + relative program arg with the `__NULLEXIT_HOME__` sentinel (install recipe gains a single `sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|"` step); 8 devref prose sites genericized to `~/.colima/...`, `~/Library/LaunchAgents/...`, `$USER`, `<USER>`.
2233 - **July 1, 2026 `a5e2a5a` (refactor: script-relative paths)** Replaced 6 hardcoded `/Users/<user>/.../nullexit/...` literals with script-relative resolution. `recover.sh`/`watcher.sh` inject `SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"`; `RECOVER="$SCRIPT_DIR/../recover.sh"` resolves without launch-path dependency.
2234 - **June 28, 2026 `f8f8e4e` (M1: scripts/ move)** 8 internal scripts consolidated into `scripts/`; the 4 user-facing/orchestrator/config files (`toggle.sh`, `recover.sh`, `setup.sh`, `docker-compose.yml`) stayed at repo root.
2235 - **June 28, 2026 `0875ef3` (ci: glob)** CI `py_compile` switched to `python3 -m py_compile scripts/*.py` after M1 left root with zero `.py` files.
2236 - **June 28, 2026 `25b37a1` (docs: scripts/ prefix)** `devref.md` + `README.md` updated to the `scripts/` prefix for all 8 moved files (~31 patches).
2237
2238 ### End-to-end verification (commit `c5b92c0`)
2239 Manual START STOP cycle ran clean on macOS after path relativization + personal-leak cleanup:
2240 - `bash toggle.sh` START: exit 0, ~135s (cold Colima boot); marker present (`2026-07-02T03:41:41Z`), all 5 containers healthy, host DNS hijacked to `100.100.21.8` (no fallback), exit node selected, Cloudflare trace through WARP succeeds.
2241 - `bash toggle.sh` STOP: exit 0, ~12s; marker cleared, all nullexit containers gone, host DNS restored to `1.1.1.1`, Tailscale disconnected.
2242
2243 ## 18. TODO & Future Work
2244
2245 ### 18.1 TODO
2246 * **FaceTime / real-time P2P split-route (Option 2, narrow):** ~~Implement `add_apple_split_routes()` / `remove_apple_split_routes()` in `common.sh` (route FaceTime's Apple `/16`s direct via the physical gateway, opt-in `.env` flag, default OFF), wired into toggle START/STOP + `recover.sh --post-wake`. **Blocked on** an empirical `tcpdump` capture of the relay/STUN `/16`s during a real call (user-triggered).~~ **Closed** implemented and shipped enabled by default. The `tcpdump` capture from a real FaceTime call confirmed `17.249.0.0/16` as the load-bearing media/relay range; `FACETIME_DIRECT_SUBNETS` in `.env.example` ships that value. Full design + diagnosis in 15.12.22.
2247 * **Fix the shadowed exit-path port REJECTs (DoT/QUIC bypass):** ~~The `FORWARD` `-i tailscale0 -p udp --dport 853`/`--dport 443` REJECT rules are shadowed by the earlier `tailscale0  tun0` ACCEPT (0 pkts), so the intended block-DoT / force-AdGuard and QUIC-block are **not enforced** a mesh client can bypass AdGuard via DoT (853). Re-order the REJECTs before the ACCEPT (or scope `-o tun0`).~~ **Closed** moved the REJECT rules in `post-rules.txt` before the `tailscale0  tun0` / `tun0  tailscale0` ACCEPT rules. Live `iptables -L FORWARD -n -v --line-numbers` now shows the REJECTs ahead of the ACCEPTs and actually accumulating hits. See 15.12.22.
2248 * **Decide toggle STOP/START from persisted *intent*, not a live probe:** ~~`toggle.sh` decided STOP-vs-START by calling `is_gateway_active()`, which live-probes the running system (`docker compose ps` + DNS + SOCKS checks). When Colima/Docker was down the `docker compose ps` blocked for the full 15 s `run_with_timeout` window and, under `set -e` + the `ERR` trap, aborted the script *before the START body ran* so the path that would boot Colima never executed (`EXIT: code=1 phase=start`, seen 57 historically).~~ **Closed** added `gateway_should_be_running()` in `common.sh`, which reads persisted intent (`MARKER_FILE`) captured into `GATEWAY_MARKER_AT_STARTUP` before the defensive stale-marker clear, and only falls back to `is_gateway_active` when the marker is absent *and* the Docker socket is actually reachable (so a dead-socket probe can neither hang nor trip `set -e`). All four decision sites (`toggle.sh` main + `--restart`, `toggle-linux.sh` main + `--restart`) now use it, and the STOP `docker compose down` is `|| true`-guarded so a disable always completes even with Docker down. See 15.12 Misc Resolved.
2249 * **Verify Wi-Fi Roaming:** Investigate and test whether switching Wi-Fi networks now successfully and smoothly recovers the gateway end-to-end without any tailscale flag errors or dead tunnels. (Pending user verification).
2250 * **Fix Honey-Port subshell fd inheritance (15.12.24):** The tripwire's arming subshell (`toggle.sh` ~L1704) only redirects the inner `nc -l` command's stdout/stderr, not its own so it holds any caller's pipe open forever once armed. Redirect the subshell itself (`(  ) >/dev/null 2>&1 &`) at both arming sites (macOS + Linux) so piping `toggle.sh`'s output no longer hangs post-completion.
2251 * **Fix Diagnostic Script Check 5/8:** ~~Investigate why `diagnose-host-leak.sh` check 5/8 (`Host default route`) sometimes reports "default route goes via physical Wi-Fi Tailscale route assertion failed" on macOS even though the `warp=on` egress checks confirm that traffic is successfully tunneling.~~ **Closed** fixed by switching check 5 from `netstat -rn | grep default` to `route -n get 1.1.1.1`. macOS Tailscale uses interface-scoped routes and does not replace the DHCP default route entry. See 15.7.1 Known Unknowns.
2252 * **Expose Tor Control Port (9051):** ~~Consider mapping the Tor Control port to localhost and adding `nyx` (the Tor terminal status monitor) to allow users to inspect their entry, middle, and exit node IPs in real-time.~~ **Closed** mapped the Tor Control Port to a procedurally generated, localhost-bound port `TOR_CONTROL_PORT`, configured `ControlPort` and disabled cookie authentication in `docker/tor/entrypoint.sh`, and integrated it into the `/sweep` verification protocol.
2253 * **Host-Only Tor Mode (Pluggable Egress):** Implement an `EGRESS_TYPE=tor_host_only` mode for users in heavily censored (DPI-restricted) environments. This mode would skip booting `warp` and `tailscale`, boot `adguardhome` and `tor` (using Telegram `obfs4` bridges), set AdGuard's upstream to Tor's `DNSPort`, and use the `pf` kill-switch to force all host Mac traffic strictly through the Tor network. This provides a 100%

```
    free, DPI-resistant evasion path without requiring an external VPS.
* **Real-activity decoy traffic (upgrade to `noise.sh`):** `noise.sh`'s current cover traffic is one-directional random UDP padding toward the exit  it blurs uplink volume/timing but does nothing for downlink (the larger, more revealing direction: page loads, images, video) and can't hide any real spike bigger than its own modest rate (15.3.1's battery-driven decision not to saturate the link). The upgrade: instead of junk padding, periodically fire *real* outbound requests (page fetches, DNS lookups) that get *real* responses  this naturally covers downlink too, since a genuine response comes back at a genuine size/timing, not a synthetic one. **Not a clean win, two open problems:** (1) making the decoy request *shape* (timing, correlation between a page load and its embedded assets) statistically indistinguishable from genuine human browsing is itself the open research problem of *website fingerprinting*  published attacks already achieve real, nontrivial classification accuracy against far more mature systems than this one; a naive independent-random-fetch generator would likely be distinguishable from genuine sessions by the same techniques. (2) still costs real bandwidth to be big enough to mask a genuinely large spike  same fundamental cost as raw padding, just spent fetching real content instead of junk. Also carries a practical/etiquette concern raw padding doesn't: firing synthetic requests at real third-party servers at a "human-like" rate is closer to bot traffic than background noise, and could trigger rate-limiting or abuse detection on the destination's side. Research-grade, not a quick toggle.
  * **The Technical Realities (What You Need to Watch Out For):**
    * **The UDP Problem:** Tor only transports TCP traffic. It does not support UDP (aside from DNS requests passed directly to its DNSPort). macOS is incredibly chatty over UDP (FaceTime, mDNS, QUIC protocols). Your pf rules defined in `scripts/pf.conf` must be configured to aggressively and silently DROP all host UDP traffic (except for local DNS queries to AdGuard). If you don't drop it cleanly, macOS services might hang indefinitely while waiting for timeouts.
    * **The "Daily Driver" Friction:** Routing a whole host OS through Tor is brutal on performance. Background daemons (iCloud sync, macOS software updates, Spotlight indexing) will blindly attempt to download gigabytes of data over Tor, which could saturate the circuits or time out completely.
    * **Bridge Rot:** State firewalls actively hunt and block obfs4 bridges. When a user's bridges burn out, they will lose all internet access due to the pf kill-switch. You will need to ensure the user has a frictionless way to update the `tor-bridges.txt` file and restart the Tor container without exposing their IP.
    * **Exit Node Discrimination:** Because all traffic will exit from known Tor nodes, the user will face an onslaught of Cloudflare CAPTCHAs, and many banking or streaming services will flat-out reject the connections.

### 18.2 Future Work
- **Direct P2P on VM Hosts:** Non-Linux hosts run Docker inside a VM, making true UDP hole-punching **to arbitrary external peers** impossible; those peers fall back to the DERP relay unless bridged mode is enabled on a trusted LAN. The only fully general solution is a native Linux host (Raspberry Pi, Intel NUC) where Docker runs without VM hypervisor translation. *(The one case that works regardless is direct **hostcontainer** P2P  the gateway and its own Mac, the same physical machine  re-permitted via routing-fix `.host_ips` RETURN rules and confirmed `direct` even in plain user-mode networking; see 15.3.5 and 15.8.7.)*
- **Native Linux Deployment:** Benchmark on a Raspberry Pi  native container footprint is ~75MB without macOS hypervisor overhead.
- **Mesh-Wide Filesystem (SFTP over Tailscale):** Any mesh device passively exposes its filesystem over SFTP. Android via Termux + OpenSSH + wakelock; iOS is client-only due to background process limits.
- **Post-Quantum Cryptography (PQC):** Tailscale uses Curve25519, theoretically vulnerable to "harvest now, decrypt later" quantum attacks. A future iteration could replace the Tailscale container with raw WireGuard + [Rosenpass](https://rosenpass.eu/) to negotiate post-quantum PSKs.
- **Egress Compartmentalization (Pluggable Architecture):** See 12.9 for the full plan to make the entire egress layer swappable via a single `.env` variable.
```

5 diagrams.md

```
# nullexit  Flow Diagrams & Architecture

---

## 1. System Architecture  What's Running & Where

```mermaid
graph TB
    subgraph INTERNET["  Internet / Cloudflare Edge"]
        CF["Cloudflare WARP\n(Exit Point)\nwarp=on"]
    end

    subgraph TAILNET[" Tailscale Mesh (100.x.x.x)"]
        PHONE[" Phone / Laptop\n(Tailnet Client)"]
    end

    subgraph MACHOST[" macOS Host"]
        direction TB
        PFFIREWALL["macOS pf Firewall\n(Kill-Switch & TCP MSS Clamping)"]
        TSDAEMON["tailscaled\n(brew service)\nhost Tailscale"]
        HOSTDNS["Host DNS\n Gateway IP\n(hijacked by toggle.sh)"]
        HOSTSOCKS["Host SOCKS5\n127.0.0.1:1080\n(fallback only)"]
```
```

```
        HOSTTOR["Host Tor SOCKS5\n127.0.0.1:$TOR_SOCKS_PORT\n(randomized)"]
        HOSTTORCTRL["Host Tor Control\n127.0.0.1:$TOR_CONTROL_PORT\n(randomized)"]

        subgraph MONITORS[" Background Daemons (toggle.sh children)"]
            WARPWATCH["WARP Watcher\n(every 30s, via docker exec)\n output.log"]
            DNSWATCHER["DNS Watcher\n(re-hijacks if drift detected)\n output.log"]
            CAFFEINATE["caffeinate -i\n(sleep prevention)"]
        end

        subgraph COLIMAVM["  Colima VM (Linux, 600MB RAM)"]
            subgraph NETNS["warp container network namespace (shared by ALL containers)"]
                GLUETUN["warp / Gluetun\nWireGuard  tun0\n(owns the namespace)"]
                TS["tailscale container\nAdvertises exit node\ntailscale0 interface"]
                SOCKS["tailscale\nSOCKS5\nport 1080"]
                ADGUARD["adguardhome\nDNS sinkhole\nport 5335"]
                ROUTINGFIX["routing-fix sidecar\nGeo-IP Blocking & Go Logger\n(writes blocked.log)"]
                TOR["tor container\nSOCKS5: port 9050\nControl: port 9051\nTransparent Proxy: port 9040\
    nDNSPort: port 5353"]
            end
        end

        subgraph LAUNCHD[" launchd (always-on)"]
            WATCHER["scripts/watcher.sh\n(sleep/wake + Wi-Fi roam\n+ LAN P2P auto-detection\n+ Wi-Fi auto-
    rejoin  always on)"]
        end
    end

    subgraph RECOVERY[" Recovery"]
        RECOVER["recover.sh\n(nuclear reset)"]
        RECOVERPOST["recover.sh --post-wake\n(light refresh)"]
    end

    %% Traffic flow
    PHONE -->|"WireGuard / Tailscale mesh"| TSDAEMON
    TSDAEMON -->|"exit-node route utun*"| TS
    TS -->|"FORWARD  tun0 (MASQUERADE)"| GLUETUN
    GLUETUN -->|"WireGuard UDP (macOS Host Egress)"| PFFIREWALL
    PFFIREWALL -->|"TCP MSS Clamped / Kill-Switch check"| CF
    TS -->|"198.18.0.0/15 PREROUTING redirect"| TOR
    TOR -->|"internal path / exit nodes via tun0"| GLUETUN

    %% DNS
    HOSTDNS -->|"UDP:53  TCP:5354"| ADGUARD
    ADGUARD -->|"DNS upstream via tun0"| CF
    ADGUARD -->|"onion queries  localhost:5353"| TOR

    %% Monitoring
    WARPWATCH -->|"docker exec warp wget cdn-cgi/trace"| GLUETUN
    DNSWATCHER -->|"networksetup -getdnsservers"| HOSTDNS

    %% Logging
    ROUTINGFIX -->|"writes"| BLOCKEDLOG["blocked.log\n(Host-side)"]
    HOSTTOR -->|"port mapping"| TOR
    HOSTTORCTRL -->|"port mapping"| TOR

    %% Recovery triggers
    WARPWATCH -->|"6 consecutive warp=off"| RECOVER
    WATCHER -->|"sleep/wake / roam"| RECOVERPOST
    RECOVERPOST -->|"if gateway unhealthy"| RECOVER
    WATCHER -->|"writes .lan_p2p_detected"| ROUTINGFIX

    style MONITORS fill:#1a1a2e,color:#e0e0ff
    style COLIMAVM fill:#0d2137,color:#e0e0ff
    style NETNS fill:#0a1628,color:#e0e0ff
    style RECOVERY fill:#2d0a0a,color:#ffe0e0
    style INTERNET fill:#0a2d1a,color:#e0ffe0
    style TAILNET fill:#1a2d0a,color:#e8ffe0
    style LAUNCHD fill:#2d1a2d,color:#ffe0ff
```

---

2. toggle.sh Full START Flow

```mermaid
flowchart TD
    START(["./toggle.sh"])
    CHECK{"is_gateway_active?\n(containers running OR\nDNS  1.1.1.1)"}

    START --> CHECK
    CHECK -->|"YES  already ON"| STOPFLOW[" See STOP Flow"]
```

```
    CHECK -->|"NO  start it"| S1

    S1["1. Reset DNS  1.1.1.1\n(prevent deadlocks)"]
    S2["2. tailscale down\n(prevent exit-node routing during boot)"]
    S3["3. Check Colima VM\n(colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged)"]
    S4["4. Configure VM swap\n(400MB, prevent OOM)"]
    S5["5. Clean corrupted AdGuard config\n(prevent container crash loop)"]
    S6["6. docker compose up -d --build\n(compile Go threat rules & boot containers)"]
    S7["7. Poll tailscale status\n(wait for 'offers exit node'\nup to 60s)"]
    TSREADY{"container\nTailscale ready?"}
    ABORT([" ABORT\n cleanup_handler"])
    S8["8. Resolve gateway Tailscale IP\n(.gateway_ip or docker exec)"]
    S9["9. Verify tailscaled running\n(auto-start if needed)"]
    S10["10. tailscale up --reset\n--exit-node=\n--accept-routes=true\n--ssh=true --accept-dns=false"]

    PF1{"Pre-flight 1/3\ntailscale ping gateway"}
    PF2{"Pre-flight 2/3\ndig +tcp AdGuard DNS"}
    PF3{"Pre-flight 3/3\nWARP internet check"}
    ALLPASS{"All 3 pass?"}

    EXITPATH[" EXIT NODE PATH\n tailscale up --exit-node=\n force_dns_to_gateway (3 retry)\n SOCKS5 disabled"]
    SOCKSPATH[" SOCKS5 FALLBACK PATH\n macOS SOCKS5  127.0.0.1:1080\n DNS proxy  127.0.0.1:53\n No exit node (
     SKIP_EXIT_NODE=true)"]

    DAEMONS["Start background daemons\n caffeinate -i (sleep prevention)\n DNS Watcher\n WARP Watcher\n(all
     write  output.log)"]

    RULECOMPILE["Rule compiler status\n(log rule/IP counts)"]
    WRITEMARKER["write_gateway_active_marker\n(/tmp/nullexit-gateway-active.marker)"]
    DONE([" Gateway UP"])

    S1 --> S2 --> S3 --> S4 --> S5 --> S6 --> S7
    S7 --> TSREADY
    TSREADY -->|"NO (40 NoState)"| ABORT
    TSREADY -->|"YES"| S8
    S8 --> S9 --> S10
    S10 --> PF1 --> PF2 --> PF3
    PF1 & PF2 & PF3 --> ALLPASS
    ALLPASS -->|"YES"| EXITPATH
    ALLPASS -->|"NO"| SOCKSPATH
    EXITPATH --> DAEMONS
    SOCKSPATH --> DAEMONS
    DAEMONS --> RULECOMPILE --> WRITEMARKER --> DONE

    style ABORT fill:#5c1010,color:#fff
    style EXITPATH fill:#0a3d1a,color:#c0ffc0
    style SOCKSPATH fill:#3d2d00,color:#ffe0a0
    style DAEMONS fill:#0a1f3d,color:#c0d8ff
    style DONE fill:#0a3d1a,color:#fff
```

---

3. toggle.sh STOP Flow

```mermaid
flowchart TD
    STOP(["./toggle.sh\n(gateway already ON)"])

    ST1["1. tailscale down\n(disconnect from mesh)"]
    ST2["2. docker compose down -t 5\n(stop all containers)"]
    ST3["3. Reset DNS  1.1.1.1\n(networksetup on all services)"]
    ST4["4. stop_local_dns_proxy\n(kill Python UDP proxy)"]
    ST5["5. stop_pidfile_daemon\n(kill caffeinate, rm PID)"]
    ST6["6. stop_pidfile_daemon\n(kill DNS Watcher, rm PID)"]
    ST7["7. stop_warp_watcher\n(kill WARP Watcher, rm PID)"]
    ST8["8. clear_gateway_active_marker\n(rm /tmp/nullexit-gateway-active.marker)"]
    ST9["9. cleanup_network_state\n disable_all_proxies\n Flush DNS cache (dscacheutil)\n Flush stale routes\n
     Power-cycle Wi-Fi (offon)"]

    DONE([" Gateway DOWN\nInternet restored"])

    STOP --> ST1 --> ST2 --> ST3 --> ST4
    ST4 --> ST5 --> ST6 --> ST7 --> ST8 --> ST9 --> DONE

    style DONE fill:#0a3d1a,color:#fff
```

---

4. Monitoring Layer The 3 Background Daemons

```
```mermaid
graph LR
    subgraph DAEMONS["Background Daemons (all start with gateway, stop on teardown)"]
        direction TB

        subgraph D1[" Sleep Prevention"]
            CAFF["caffeinate -i\nPID: /tmp/nullexit-caffeinate.pid\n\nKeeps Mac awake while\ngateway is active
    "]
        end

        subgraph D2[" DNS Watcher"]
            DNSW["Polls networksetup every ~30s\nPID: /tmp/nullexit-dns-watcher.pid\n\nIf DNS drifts from
    gateway IP\n re-hijacks automatically\n output.log"]
        end

        subgraph D3[" WARP Watcher"]
            WARPW["Polls every 30s via:\ndocker compose exec warp wget cdn-cgi/trace\nPID: /tmp/nullexit-warp-
    watcher.pid\n\nChecks uplink first (link-state):\nno IP, or link 'inactive' (stale lease)\n PAUSE (rejoin
    Wi-Fi), not a kill\n(15.6.2 / 15.12.27)\n\nElse counts consecutive warp=off\n WARP_FAIL_THRESHOLD (default
    6=30s)\n runs recover.sh (nuclear)\n output.log"]
        end

    end

    subgraph BLIND["What each monitor can/can't see"]
        B1["WARP Watcher sees:\n Container WARP tunnel state\n Host NIC traffic\n Flash leaks < 5s"]
        B2["Host-NIC blind spot covered by:\n PF kill-switch (kernel, fail-closed)\n on-demand: diagnose-host-
    leak.sh\n(--watch to monitor) / sweep.sh"]
    end

    D3 -.->|"covers"| B1
    B1 -.->|"gap closed by"| B2

    style D1 fill:#1a2d0a
    style D2 fill:#0a1a2d
    style D3 fill:#2d0a2d
    style BLIND fill:#1a1a1a,color:#aaa
```

---

5. Traffic Flow Data Path (Normal Operation)

```mermaid
sequenceDiagram
    participant PHONE as  Phone<br/>(Tailnet client)
    participant TS_HOST as tailscaled<br/>(macOS host)
    participant TS_CONTAINER as tailscale<br/>(container)
    participant GLUETUN as Gluetun/warp<br/>(container)
    participant CF as  Cloudflare<br/>WARP Edge
    participant INTERNET as  Internet

    Note over PHONE,INTERNET: Normal EXIT NODE path (all 3 pre-flights passed)

    PHONE->>TS_HOST: WireGuard packet<br/>(encrypted, UDP 41641)
    TS_HOST->>TS_CONTAINER: route via utun*  tailscale0
    TS_CONTAINER->>GLUETUN: FORWARD chain<br/>(iptables MASQUERADE on tun0)
    GLUETUN->>TS_HOST: WireGuard tunnel (tun0) egresses macOS Host
    Note over TS_HOST: macOS pf Firewall applies<br/>Kill-Switch & TCP MSS Clamping
    TS_HOST->>CF: re-encrypted UDP packet (MSS 1160)
    CF->>INTERNET: Plain HTTPS from<br/>Cloudflare IP

    Note over PHONE,INTERNET: DNS Resolution path
    PHONE->>TS_HOST: DNS query
    TS_HOST->>TS_CONTAINER: UDP:53  TCP:5354  AdGuard :5335
    TS_CONTAINER->>GLUETUN: AdGuard upstream DNS via tun0
    GLUETUN->>CF: Encrypted DNS query
    CF-->>GLUETUN: DNS response
    GLUETUN-->>TS_CONTAINER: response
    TS_CONTAINER-->>TS_HOST: response
    TS_HOST-->>PHONE: DNS answer
```

---

6. recover.sh Decision Tree

```mermaid
flowchart TD
    INVOKE(["recover.sh called"])
```
```

```
    MODE{"--post-wake\nflag?"}

    subgraph POSTWAKE[" --post-wake (light refresh, gateway stays UP)"]
        PW1["Re-hijack DNS  gateway IP"]
        PW2["Flush DNS cache"]
        PW3["Refresh tailscale exit-node"]
        PW4["Force-recreate warp container\n(if Gluetun healthcheck failed)"]
        PW5["Leave: containers, Wi-Fi,\ncaffeinate, DNS watcher  untouched"]
    end

    subgraph NUCLEAR[" Default (nuclear  gateway torn down)"]
        N0["Sudo keep-alive loop (background)"]
        N1["1. tailscale down\n(disconnect from mesh)"]
        N2["2. Reset DNS  1.1.1.1\n(ALL network services)"]
        N3["3. disable_all_proxies\n(all interfaces)"]
        N4["4. Flush DNS cache\n(dscacheutil -flushcache)"]
        N5["5. Flush stale routes\n(WARP endpoint routes)"]
        N6a["6a. docker compose down -t 30\n(stop all containers)"]
        N6b["6b. stop_pidfile_daemon\n(kill caffeinate)"]
        N6c["6c. stop_pidfile_daemon\n(kill DNS Watcher)"]
        N7["7. Power-cycle Wi-Fi\n(networksetup offsleep 2on)"]
        N8["8. Verify internet working\n(curl ifconfig.io)"]
    end

    DONE_PW([" Gateway refreshed\n(still running)"])
    DONE_NUC([" Internet restored\n(gateway fully stopped)"])

    INVOKE --> MODE
    MODE -->|"yes"| PW1
    PW1 --> PW2 --> PW3 --> PW4 --> PW5 --> DONE_PW
    MODE -->|"no"| N0
    N0 --> N1 --> N2 --> N3 --> N4 --> N5
    N5 --> N6a --> N6b --> N6c --> N7 --> N8 --> DONE_NUC

    style POSTWAKE fill:#0a2d1a,color:#c0ffc0
    style NUCLEAR fill:#2d0505,color:#ffc0c0
    style DONE_PW fill:#0a3d1a,color:#fff
    style DONE_NUC fill:#3d1a00,color:#ffe0c0
```

---

7. Failure Paths & Self-Healing

```mermaid
flowchart LR
    subgraph EVENTS["Failure / Event"]
        E1["WARP tunnel drops\n(warp=off)"]
        E2["Mac wakes from sleep\nor Wi-Fi roams"]
        E3["DNS drifts from\ngateway IP"]
        E4["toggle.sh crashes /\nCtrl-C / SIGTERM"]
        E5["Host-side leak\n(non-tunnel HOST NIC egress)"]
        E6["Silent NAT timeout\n(ping stops working)"]
        E7["Wi-Fi uplink drops\n(auto-join off  no internet)"]
    end

    subgraph DETECTORS["Detector"]
        D1["WARP Watcher\n(30s poll, docker exec)"]
        D2["scripts/watcher.sh\n(launchd, always on)"]
        D3["DNS Watcher\n(30s poll, networksetup)"]
        D4["cleanup_handler\n(ERR/INT/TERM/HUP trap)"]
        D5["PF kill-switch\n(kernel, always-on) +\ndiagnose-host-leak.sh --watch\n(on-demand)"]
        D6["routing-fix.sh\n(30s curl over tun0)"]
    end

    subgraph ACTIONS["Response"]
        A1["uplink up +  threshold off\n recover.sh (nuclear).\nuplink DOWN  PAUSE, not kill"]
        A2["recover.sh --post-wake\n(gentle refresh)"]
        A3["Re-hijack DNS\n(networksetup setdnsservers)"]
        A4["Stop all daemons\nReset DNS  1.1.1.1\nTear down Tailscale"]
        A5["Non-tunnel egress blocked\nfail-closed at kernel;\n--watch alerts on state change"]
        A6["Blackhole internet traffic\n+ Drop TUNNEL_FAILED_CLOSED.marker"]
        A7["Push macOS Desktop Notification\n(via watcher.sh listener)"]
        A8["watcher.sh Listener 4 \nwifi-rejoin.sh: re-associate a\nremembered network (grace+backoff)"]
    end

    E1 --> D1 --> A1
    E2 --> D2 --> A2
    E3 --> D3 --> A3
    E4 --> D4 --> A4
    E5 --> D5 --> A5
```

```
    E6 --> D6 --> A6
    E7 --> D2 --> A8
    A6 -.->|"marker detected"| D2
    D2 --> A7

    style EVENTS fill:#2d1010,color:#ffcccc
    style DETECTORS fill:#0a1a2d,color:#cce0ff
    style ACTIONS fill:#0a2d0a,color:#ccffcc
```

---

8. output.log Who Writes What

All events converge into a single `output.log` at the repo root.

Writer	Event Types	Format
`toggle.sh`	All startup/shutdown steps, errors, pre-flight results	Plain text with timestamps
`recover.sh`	Every recovery step (nuclear or post-wake)	Plain text
WARP Watcher	`WARP DOWN`, `WARP RECOVERED`, `WARP SHUTDOWN`	`[UTC timestamp]` prefix
DNS Watcher	DNS re-hijack events	Appended inline
`docker compose`	Container logs on failure (last 100 lines of warp)	Dumped on ERR path
`routing-fix.sh`	(via `nullexit-logger` inside container)	Structured

Grep cheatsheet:
```bash
grep 'WARP DOWN' output.log    # Container-side tunnel drops
grep 'WARP SHUTDOWN' output.log # Auto-recovery triggered
grep 'WARP RECOVERED' output.log # Self-healed before threshold
grep -E 'EXEC:|EXIT ' output.log # Lifecycle breadcrumbs (last command + exit)
```

6 toggle.sh

```bash
#!/bin/bash
# toggle.sh  nullexit gateway toggle script (macOS + Linux)
# Automatically handles Docker containers, Colima, DNS hijacking, and Tailscale exit node routing.
#
# One file, both OSes: the dispatch below runs the self-contained Linux
# implementation (shuf / ss / nmcli / ip / systemd) and exits; on macOS it
# falls through to the macOS implementation (jot / netstat / networksetup,
# launchd / pf).

set -e

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"

#
# LINUX TOGGLE  self-contained; runs the full toggle then exits before the
# macOS body below. Native Linux tooling (shuf / ss / nmcli / ip / systemd).
# (Body inlined verbatim at column 0  it contains multiline `bash -c ""`
# strings that must NOT be re-indented.)
#
if [[ "$OSTYPE" != "darwin"* ]]; then
# Define log file (used by the lifecycle breadcrumb + run_logged helper)
LOG_FILE="$PWD/output.log"

# --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
# TOGGLE_PHASE is updated as the script moves through STOP/START so that if the
# process is killed (terminal closed, pkill, OOM) or exits on error, the EXIT
# trap records the final exit code AND the phase it died in  making an
# interrupted START traceable in output.log instead of ending abruptly.
TOGGLE_PHASE="init"
_log_exit_breadcrumb() {
  local rc=$?
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
    2>/dev/null || true
}
trap '_log_exit_breadcrumb' EXIT

# --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
# When enabled, mirror a full command-by-command xtrace to output.log. Off by
# default. Read directly from .env because common.sh (read_env_var) isn't
# sourced yet. Branch on bash version: BASH_XTRACEFD needs >= 4.1 (Linux
# usually has bash 5); fall back to teeing stderr on older bash. We never use
```

```
# the `exec {var}>>` dynamic-fd form (4.1+ only).
DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'" | tr '[:upper:]' '[:lower:]' | tr -d '[:space:]')
if [ "$DEBUG_TRACE" = "true" ]; then
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled  xtrace mirrored to output.log (bash ${BASH_VERSINFO[0]}.${BASH_VERSINFO[1]})" >> "$LOG_FILE"
  export PS4='+ [${SECONDS}s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: '  # subshell-free clock: a $() in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
  if [ "${BASH_VERSINFO[0]}" -gt 4 ] || { [ "${BASH_VERSINFO[0]}" -eq 4 ] && [ "${BASH_VERSINFO[1]}" -ge 1 ]; }; then
    exec 9>>"$LOG_FILE"
    export BASH_XTRACEFD=9
    set -x
  else
    # bash < 4.1: plain stderrfile redirect. NOT a process substitution  under
    # set -e, xtrace flowing through a backgrounded `tee` can abort the script
    # (failed write/SIGPIPE trips set -e). See devref 15.11.8.
    exec 2>>"$LOG_FILE"
    set -x
  fi
fi

# --- PROCEDURAL PORT GENERATION ---
# To prevent static port fingerprinting by malware, we randomize exposed ports
# on every boot and persist them to a hidden environment file.
PORTS_FILE="/tmp/nullexit-ports.env"
if [[ "$1" == "" || "$1" == "--restart" || "$1" == "restart" ]]; then

  # Collision-avoidance function
  GENERATED_PORTS=""
  get_free_port() {
    local port
    while true; do
      # Linux uses shuf or $RANDOM instead of jot
      port=$(shuf -i 10000-65000 -n 1 2>/dev/null || echo $((10000 + RANDOM % 55000)))
      # 1. Prevent self-collision within this loop
      if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
        continue
      fi
      # 2. Prevent collision with ANY process on the machine (root daemons, terminal apps, etc)
      # We use ss to read the kernel socket table directly.
      if ! ss -lnt | awk '{print $4}' | grep -q ":$port$"; then
        echo "$port"
        return
      fi
    done
  }

  export TOR_SOCKS_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_SOCKS_PORT"

  export TOR_CONTROL_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"

  export SOCKS_PROXY_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"

  export DNS_PROXY_PORT=$(get_free_port)

  echo "export TOR_SOCKS_PORT=$TOR_SOCKS_PORT" > "$PORTS_FILE"
  echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
  echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
  echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
elif [ -f "$PORTS_FILE" ]; then
  # On STOP, source the existing ports so docker-compose down matches
  source "$PORTS_FILE"
else
  # Fallbacks if file is lost
  export TOR_SOCKS_PORT=9050
  export TOR_CONTROL_PORT=9051
  export SOCKS_PROXY_PORT=1080
  export DNS_PROXY_PORT=5354
fi

# Start execution timer
TOGGLE_START_TIME=$SECONDS

if [[ "$1" == "restart" ]] || [[ "$1" == "--restart" ]]; then
  echo "Executing Gateway Restart Sequence..."
  # We must source common.sh here temporarily to check gateway state before we proceed
  source "$SCRIPT_DIR/scripts/common.sh"
```

```
  # Use persisted intent via gateway_state (echoes a string, never a return code;
  # set -e-safe under set -x  see devref 15.11.8). On Linux the marker is not
  # maintained, so this degrades to the is_gateway_active fallback (fast  native Docker).
  if [ "$(gateway_state)" = "running" ]; then
    echo "Gateway is currently running. Stopping it first..."
    bash "$0"
    echo "Gateway stopped. Now starting it..."
  else
    echo "Gateway is already stopped. Starting it..."
  fi
  bash "$0"
  exit 0
fi

# Global flag to track if the script completed successfully
SUCCESS_RUN=false

# Global variable to track the currently running background command PID
CURRENT_BG_PID=""

# Source common bash functions
source "$SCRIPT_DIR/scripts/common.sh"


PID_FILE="/tmp/nullexit-caffeinate.pid"

# Start macOS caffeinate to prevent sleep while gateway is running
start_sleep_prevention() {
  if ! command -v systemd-inhibit >/dev/null 2>&1; then
    echo "  [Warning] systemd-inhibit tool not found. Sleep prevention unavailable."
    return 1
  fi

  # Stop any existing sleep prevention process first to avoid duplicates
  stop_sleep_prevention

  echo "  Enabling system sleep prevention (systemd-inhibit)..."
  # Run systemd-inhibit wrapped in a bash trap to prevent idle system sleep.
  # Critically, this traps shutdown signals (SIGTERM) and automatically
  # flushes hijacked DNS back to normal right before the PC powers off.
  # We do not use recover.sh here because it is a heavy recovery script
  # which takes too long and gets terminated before it can reset the DNS. Instead,
  # we surgically and instantly restore normal DNS settings here.
  nohup bash -c "
    trap '
      echo \"[Shutdown Trap] Reverting network settings...\" >> \"$PWD/output.log\" 2>&1
      kill \$! 2>/dev/null || true
      sudo -n resolvectl domain \"$ACTIVE_SERVICE\" \"\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n resolvectl revert \"$ACTIVE_SERVICE\" >> \"$PWD/output.log\" 2>&1 || true
      resolvectl domain \"$ACTIVE_SERVICE\" \"\" >> \"$PWD/output.log\" 2>&1 || true
      resolvectl revert \"$ACTIVE_SERVICE\" >> \"$PWD/output.log\" 2>&1 || true
      if command -v tailscale >/dev/null 2>&1; then
        tailscale up --accept-dns=false --exit-node= >> \"$PWD/output.log\" 2>&1 &
        TS_DOWN_ARGS=\"\"
        if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
        tailscale down \$TS_DOWN_ARGS >> \"$PWD/output.log\" 2>&1 &
      fi
      sleep 1
      echo \"[Shutdown Trap] Cleanup complete.\" >> \"$PWD/output.log\" 2>&1
      exit 0
    ' SIGTERM SIGINT SIGHUP
    systemd-inhibit --what=sleep --who=nullexit --why=\"gateway running\" --mode=block sleep infinity &
    wait \$!
  " >> output.log 2>&1 &
  local caffe_pid=$!
  echo "$caffe_pid" > "$PID_FILE"
  echo "  Sleep prevention active (PID $caffe_pid). Your PC won't sleep while the gateway is running."
}

# Stop sleep prevention when gateway is stopped
stop_sleep_prevention() {
  if [ -f "$PID_FILE" ]; then
    local caffe_pid
    caffe_pid=$(cat "$PID_FILE")
    if [ -n "$caffe_pid" ] && kill -0 "$caffe_pid" 2>/dev/null && ps -p "$caffe_pid" -o command= 2>/dev/null |
     grep -q -E "systemd-inhibit|recover.sh"; then
      echo "  Stopping system sleep prevention (systemd-inhibit wrapper PID $caffe_pid)..."
      kill "$caffe_pid" 2>/dev/null || true
    fi
    rm -f "$PID_FILE"
  fi
```

```
}

DNS_WATCHER_PID_FILE="/tmp/nullexit-dns-watcher.pid"

stop_dns_watcher() {
  if [ -f "$DNS_WATCHER_PID_FILE" ]; then
    local watcher_pid
    watcher_pid=$(cat "$DNS_WATCHER_PID_FILE")
    if [ -n "$watcher_pid" ] && kill -0 "$watcher_pid" 2>/dev/null; then
      echo "  Stopping background DNS Watcher (PID $watcher_pid)..."
      kill -9 "$watcher_pid" 2>/dev/null || true
    fi
    rm -f "$DNS_WATCHER_PID_FILE"
  fi
}

# Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
cleanup_handler() {
  local exit_code=$?
  local trigger_type="$1"
  local line_no="$2"

  if [ "$SUCCESS_RUN" = "false" ]; then
    echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1...
    "
    if [ -n "$ACTIVE_SERVICE" ]; then
      reset_dns
    fi

    if [ -n "$CURRENT_BG_PID" ]; then
      echo "Terminating active command (PID $CURRENT_BG_PID)..."
      kill -15 "$CURRENT_BG_PID" 2>> output.log || true
    fi

    # Disconnect host Tailscale and close the GUI application on failure
    disconnect_tailscale_host

    # Stop local DNS proxy if running
    stop_local_dns_proxy

    # Stop sleep prevention
    stop_sleep_prevention
    stop_dns_watcher

    # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
    if [ -n "$ACTIVE_SERVICE" ]; then
      cleanup_network_state
    fi


    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n=============================================="
      echo "ERROR: Script failed at line $line_no with exit code $exit_code."
      echo "=============================================="
      echo "Troubleshooting steps:"
      echo "1. Run 'colima status' to verify the VM is active."
      echo "2. Run 'docker compose ps' to check container health."
      echo "3. Run 'docker compose logs' to view application logs."
      echo "4. Run 'tailscale status' to check host Tailscale status."
      echo "=============================================="
    fi
  fi
}

trap 'cleanup_handler ERR $LINENO' ERR
trap 'cleanup_handler INT' INT
trap 'cleanup_handler TERM' TERM
trap 'cleanup_handler HUP' HUP

#  NOPASSWD sudo
# The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
# needed (resolvectl, ip, systemctl, systemd-inhibit, pkill, and
# python3 scripts/dns-proxy.py for the fallback DNS proxy).
# All privileged calls below use `sudo -n` which will run without prompting.
# No credential caching loop is needed.

# Add Homebrew and standard paths
export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"

# Check for local DNS proxy tools (used when tailnet data plane is unavailable)
# Uses a Python DNS proxy (handles TCP wire format correctly).
```

```
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_BIN=""
if command -v python3 >/dev/null 2>&1; then
  DNS_PROXY_BIN="python3 -I"
elif command -v python >/dev/null 2>&1; then
  DNS_PROXY_BIN="python -I"
fi

# Path to the DNS proxy script (sibling of this script)
DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/dns-proxy.py"

# Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
# (`brew install tailscale` + `systemctl start tailscaled`)  NOT the
# Tailscale.app GUI  so there is no .app bundle or Network Extension sandbox
# to launch, no menu-bar click to perform, and no scutil Network Service to
# start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
# if you ran the install with sudo) and the control socket is always available.
TS_BIN=""
if command -v tailscale >> output.log 2>&1; then
  TS_BIN="tailscale"
fi

# Baseline: disable Tailscale DNS management at the daemon level.
# `tailscale set` writes a persistent preference. However, `tailscale up --reset`
# (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
# re-enable `--accept-dns=true` and nullify this `set`.
# Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
# `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
#
# This initial `set` is still useful as a baseline for non-reset operations
# (e.g. if the user runs `tailscale up` manually without --reset) and covers
# the brief window between this line and the first `--reset` call.
#
# Runs once at script start while tailscaled is still connected (if it was
# left running from a previous toggle), so the pref takes effect before any
# disconnection or reconnection logic.
if [ -n "$TS_BIN" ]; then
  $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
fi

# Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
# *entire* teardown  no need to mess with system network managers.
# tailscaled keeps running (as a systemd service), which is intentional:
# the next `tailscale up` then reconnects in seconds rather than waiting for a
# slow daemon launch. The `tailscale down` call also retains the user's
# auth state, so we don't force a browser re-login every cycle.
#
# Defined early (before the if/else branches and referenced by cleanup_handler's
# ERR trap) so that any failure before the main logic can still tear down Tailscale.
restart_tailscaled_daemon() {
  echo "  [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
  echo "  [Tailscale Recovery] Attempting to restart tailscaled service via systemctl..."
  if run_with_timeout 15 sudo -n systemctl restart tailscaled >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled service."
  else
    echo "  [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
  fi
  sleep 3
}
disconnect_tailscale_host() {
  if [ -n "$TS_BIN" ]; then
    echo "Disconnecting host Tailscale from mesh (tailscaled stays running as system service)..."
    local ts_args=""
    if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
    if ! run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1; then
      restart_tailscaled_daemon
      run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1 || true
    fi
  fi
}

ACTIVE_SERVICE=$(get_active_service)

# Helper to nuke leftover network state after teardown (proxy settings, routing
# table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
# problem where stale state kills internet after the gateway stops.
cleanup_network_state() {
  echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."

  # 1. Disable any leftover proxy settings (needs root on many macOS versions)
  # (Skipped for Linux since we don't set global SOCKS proxy)
```

```
  echo "  Proxies disabled."

  # 2. Flush DNS cache
  if command -v resolvectl >> output.log 2>&1; then
    sudo -n resolvectl flush-caches 2>> output.log || true
  elif command -v systemd-resolve >> output.log 2>&1; then
    sudo -n systemd-resolve --flush-caches 2>> output.log || true
  fi
  echo "  DNS cache flushed."

  # 3. Flush stale routing table entries
  if command -v ip >> output.log 2>&1; then
    sudo -n ip route flush cache >> output.log 2>&1 || true
  fi
  echo "  Routing table flushed."

  # 4. Power-cycle Wi-Fi to clear any lingering interface state
  echo "  Restarting network interface..."
  sudo -n ip link set dev "$ACTIVE_SERVICE" down 2>> output.log || true
  sleep 1
  sudo -n ip link set dev "$ACTIVE_SERVICE" up 2>> output.log || true
  echo "  Interface power-cycled ($ACTIVE_SERVICE)."
  sleep 3

  echo "Network state cleanup complete."
}

# Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
# during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
# `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
# would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
# Capture current DNS before resetting so we can check if it was hijacked
INITIAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head -1 || true)

echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
reset_dns


  #  Local Network Surveillance Check
  echo -e "\n"
  echo "Local Network Surveillance Check"
  echo ""
  # Count populated ARP cache entries to detect network scanning or high congestion
  if command -v ip >/dev/null 2>&1; then
    ARP_COUNT=$(ip neigh | grep -iv 'incomplete' | wc -l | tr -d ' ')
  else
    # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
    ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
  fi
  if [ "$ARP_COUNT" -gt 15 ]; then
    echo -e "  \033[1;33m[!] WARNING: High ARP activity detected ($ARP_COUNT devices in cache).\033[0m"
    echo -e "  \033[1;33m[!] This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan).\033[0m"
    echo -e "  [] Your traffic remains fully encrypted and invisible.\n"
  else
    echo -e "  [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
  fi

echo "Checking Gateway Status..."
# Decide STOP vs START from persisted intent via gateway_state, which echoes a
# string (never a return code) to stay set -e-safe under set -x. See devref 15.11.8.
# On Linux this falls back to is_gateway_active (fast  native Docker).
if [ "$(gateway_state)" = "running" ]; then
  TOGGLE_PHASE="stop"
  echo -e "\n=============================================="
  echo "Gateway is RUNNING. Stopping it now..."
  echo -e "==============================================\n"

  # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
  disconnect_tailscale_host
  echo ""

  echo "Stopping Docker containers..."
  # `|| true`: a disable must always complete even if the Docker daemon is down.
  run_logged docker compose down -t 5 || true

  echo -e "\nDocker daemon handles container teardown automatically."

  # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
  # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
```

```
  # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
  echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
  reset_dns

  # 3. Stop local DNS proxy if running
  stop_local_dns_proxy

  # Stop sleep prevention
  stop_sleep_prevention
  stop_dns_watcher

  # 4. Nuke leftover network state so internet actually works after teardown
  cleanup_network_state

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
else
  TOGGLE_PHASE="start"
  echo -e "\n=============================================="
  echo "Gateway is STOPPED. Starting it now..."
  echo -e "==============================================\n"

  # 1. Prevent host exit-node deadlock during VM / Container startup
  disconnect_tailscale_host

  echo -e "\nUsing native Linux Docker..."

  # 4. Clean up corrupted AdGuardHome configurations
  if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
    rm -f "adguard/conf/AdGuardHome.yaml"
    echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
  fi



  # 5. Start compose services
  write_host_ips
  # Procedurally generated ports have already been exported at the top of the script.
  echo "  [Security] Procedurally randomized proxy ports generated to evade fingerprinting."

  # --- HONEY-PORT TRIPWIRE ---
  # Generate a random trap port. If any malware does a sequential port scan on localhost,
  # it will inevitably hit this port. When it does, we instantly fire a desktop alert.
  # Reap the honey-port from any previous toggle before arming a new one, so the
  # `nc -l` tripwire never accumulates one idle listener per toggle-on (15.12.20).
  pkill -f "nc -l .*127.0.0.1" 2>/dev/null || true
  pkill -f "nc -l localhost" 2>/dev/null || true
  HONEY_PORT=$(get_free_port)
  (
    if nc -l -p "$HONEY_PORT" 127.0.0.1 >/dev/null 2>&1 || nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$PWD/output.log"
      if command -v notify-send >/dev/null 2>&1; then
        notify-send -u critical "nullexit OPSEC Alert" "An unknown application is aggressively scanning your local ports. Malware port-scan suspected."
      fi
    fi
  ) &
  disown %1 2>/dev/null || true
  echo "  [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."

  echo -e "\nStarting Docker containers..."
  run_logged docker compose up -d

  # Clean up the one-off rule compiler container so it doesn't clutter the Docker UI
  docker compose rm -s -f rule-compiler >> output.log 2>&1

  # 6. Wait for the gateway container's Tailscale connection to be ready
  BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
  USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')

  echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
  echo "  (This can take 30-60s  Tailscale goes through: NoState  Starting  Running  Online)"
  connected=false
  consecutive_failures=0
  for i in {1..60}; do
    # Run tailscale status and capture its output so the user can see what's happening.
    # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
    ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)

    if echo "$ts_output" | grep -q "offers exit node"; then
```

```
      connected=true
      break
    fi

    # Track consecutive failures to detect a stuck state.
    # If we see 40+ consecutive 'NoState' responses, give up early
    # the daemon is alive but can't connect to the coordination server.
    if echo "$ts_output" | grep -q "NoState"; then
      consecutive_failures=$((consecutive_failures + 1))
      if [ "$consecutive_failures" -ge 40 ]; then
        echo ""
        echo "  [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
        echo "    Tailscale daemon is running but can't reach the coordination server."
        break
      fi
    else
      consecutive_failures=0
    fi

    # Print a condensed status every 10 seconds so the user can see progress
    if [ $((i % 10)) -eq 0 ]; then
      ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
      echo "  [attempt $i/60] $ts_short"
    elif [ $((i % 5)) -eq 0 ]; then
      echo -n "[$i]"
    else
      echo -n "."
    fi
    sleep 1
  done
  echo ""

  if [ "$connected" = "true" ]; then
    echo "Gateway container is online on the Tailscale mesh."
  elif [ "$BYPASS_PING" = "true" ]; then
    echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
  else
    echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
    echo "  The container was seen in the following states during the wait:" >&2
    echo "$ts_output" | sed 's/^/    /' >&2
    echo "" >&2
    echo "  This could mean:" >&2
    echo "    1. Tailscale auth key has expired  check TS_AUTHKEY in .env" >&2
    echo "    2. Network issue inside the container  check: docker compose logs tailscale" >&2
    echo "    3. gluetun VPN tunnel not connecting  check: docker compose logs warp | tail -20" >&2
    echo "" >&2
    echo "  Diagnose manually:" >&2
    echo "    docker compose exec -T tailscale tailscale status" >&2
    echo "    docker compose exec -T tailscale tailscale netcheck" >&2
    cleanup_handler ERR $LINENO
    exit 1
  fi

  # 6b. Resolve the gateway's Tailscale IP.
  # Dynamically query the container for the IP (handles IP changes after re-auth)
  # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
  #
  # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
  # returns (\r) into captured output. If $TS_IP ends with \r, then
  # `networksetup -setdnsservers` silently rejects the malformed IP and
  # DNS stays at 1.1.1.1  the primary bug this project was hitting.
  # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
  # the first field of the first line.
  TS_IP=""
  if [ "$connected" = "true" ]; then
    TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
     awk 'NR==1{print $1; exit}' || true)
    if [ -n "$TS_IP" ]; then
      echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
      echo "$TS_IP" > .gateway_ip
    else
      echo "    [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
      TS_IP=$(read_adguard_ip || true)
    fi
  else
    TS_IP=$(read_adguard_ip || true)
  fi

  # 7. Connect host Mac to Tailscale mesh.
  # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
  # system LaunchDaemon via 'systemctl start tailscaled'), the previous five-
```

```
  # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
  # no menu bar item to click, and no Accessibility permission required.
  #
  # CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
  # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
  # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
  # are impossible without the host on the mesh).
  #
  # We connect in two phases:
  #   Phase A  Join mesh WITHOUT exit node (so pre-flight checks can run)
  #   Phase B  Set exit node only after pre-flight checks confirm the gateway works
  HOST_ON_MESH=false
  SKIP_EXIT_NODE=false
  if [ -n "$TS_BIN" ]; then
    echo -n "Verifying tailscaled is reachable"
    daemon_ready=false
    for i in {1..10}; do
      # `tailscale status` returns exit code 1 when the daemon is alive but
      # disconnected ("Tailscale is stopped."). We check the daemon socket
      # directly as a fallback  if the socket exists, tailscaled is running
      # even if the host isn't connected to the mesh yet.
      if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
        daemon_ready=true
        break
      fi
      echo -n "."
      sleep 1
    done
    echo ""

    if [ "$daemon_ready" != "true" ]; then
      echo -e "\033[0;33m[!] tailscaled daemon is not running. Attempting to start it...\033[0m"

      # Try system-wide daemon (with timeout  sudo may prompt for password)
      if [ "$daemon_ready" != "true" ]; then
        if run_with_timeout 10 sudo systemctl start tailscaled 2>> output.log; then
          echo "  Started tailscaled as system LaunchDaemon."
          for i in {1..15}; do
            if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
              daemon_ready=true
              break
            fi
            sleep 1
          done
        fi
      fi
    fi

    if [ "$daemon_ready" = "true" ]; then
      echo "tailscaled is responsive (running as a system service)."

      #  Phase A: Join mesh without exit node
      echo "Connecting host to Tailscale mesh (no exit node yet)..."
      if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node=; then
        HOST_ON_MESH=true
        echo "Host is on Tailscale mesh."
      else
        echo -e "\033[0;33m[!] tailscale up failed even though the daemon is running.\033[0m"
        echo "  This usually means the host hasn't authenticated with your tailnet yet."
        echo "  Run this command in a terminal to authenticate:"
        echo ""
        echo "      sudo tailscale up"
        echo ""
        echo "  A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
      fi
    else
      echo -e "\033[0;31m[!] tailscaled could NOT be started automatically.\033[0m"
      echo "  Run this in a terminal to start and authenticate Tailscale:"
      echo ""
      echo "      sudo systemctl start tailscaled"
      echo "      sudo tailscale up"
      echo ""
      echo "  Then re-run toggle.sh."
    fi
  else
    echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
  fi

  #  Phase B: Pre-flight checks + enable exit node only if safe
  if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
    echo -e "\n"
```

```
    echo "Pre-flight connectivity check"
    echo ""

    # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
    # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
    # ("direct connection not established") even though the pong was received.
    # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
    echo -n "  [1/3] Gateway reachable via Tailscale... "
    ping_ok=false
    # The host tailscaled needs a few seconds to propagate the new network map and
    # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
    for attempt in {1..45}; do
      if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
        ping_ok=true
        break
      fi
      sleep 1
    done
    if [ "$ping_ok" = "true" ]; then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
    echo -n "  [2/3] AdGuard DNS via localhost... "
    if command -v dig >/dev/null 2>&1; then
      if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 >/dev/null 2>&1; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    elif command -v nslookup >/dev/null 2>&1; then
      if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 >/dev/null 2>&1; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    else
      echo "SKIP (dig/nslookup not available)"
    fi

    # Check 3: Does the WARP container have external internet?
    echo -n "  [3/3] WARP container internet... "
    if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace &>/dev/null;
     then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Act based on check results
    if [ "$SKIP_EXIT_NODE" != "true" ]; then
      echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
      if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node="$TS_IP" --exit-node-allow-lan-access=
     true; then
        echo "Exit node enabled."
      else
        echo "[Warning] Failed to set exit node."
        SKIP_EXIT_NODE=true
      fi
    fi

    if [ "$SKIP_EXIT_NODE" = "true" ]; then
      echo -e "\n\033[0;33m[!] Pre-flight checks failed  EXIT NODE + DNS HIJACK SKIPPED.\033[0m"
      echo "  Host stays on Tailscale mesh but your internet routes through your normal connection."
      echo "  Troubleshoot with:"
      echo "    docker compose logs warp | tail -20"
      echo "    docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"
      echo ""
      echo "  Falling back to SOCKS5 proxy for traffic routing..."
    fi
  elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
    echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
    SKIP_EXIT_NODE=true
  fi

```

```
  # 8. Apply DNS Hijacking.
  # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
  # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
  # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
  if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
    echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
    echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
    if force_dns_to_gateway "$TS_IP"; then
      echo "DNS hijack verified: single server = $TS_IP."
      start_dns_watcher "$TS_IP"
    else
      echo "[Warning] DNS hijack could not be verified."
      echo "  If DNS is broken, run manually:"
      echo "    networksetup -setdnsservers \"$ACTIVE_SERVICE\" $TS_IP"
      echo "    networksetup -setsearchdomains \"$ACTIVE_SERVICE\" ts.net"
    fi
  elif [ -n "$TS_IP" ]; then
    echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
    echo "  Trying local DNS proxy for ad-blocking..."
    if start_local_dns_proxy; then
      echo "  Ad-blocking active via local DNS proxy through AdGuard."
    else
      echo "  Local DNS proxy unavailable. DNS stays at 1.1.1.1."
      echo "  AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
    fi

    # Enable SOCKS5 proxy for traffic routing through WARP
    # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
    echo -e "\n  Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
    echo "  (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
    enable_socks_proxy "$ACTIVE_SERVICE"

    # Verify the SOCKS5 proxy works through WARP
    echo -n "  Verifying SOCKS5 proxy routes through WARP... "
    if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn-cgi/
     trace 2>/dev/null | grep -q "warp=on"; then
      echo "ok (warp=on)."
      echo "  All TCP traffic now encrypted through Cloudflare WARP tunnel."
    else
      echo "FAILED (proxy not routing through WARP)."
      echo "  Disabling SOCKS5 proxy  internet stays on direct connection."
      disable_socks_proxy
      echo "  SOCKS proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
    fi
  else
    echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
  fi

  # 9. Verify host connectivity
  if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
    if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
      echo "Host is online on the Tailscale mesh."
    else
      echo "[Warning] Host is not responding on Tailscale."
    fi
  fi

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."

  # Prevent system from going to sleep while gateway is running
  start_sleep_prevention
fi

SUCCESS_RUN=true

# Final DNS state summary
if [ -n "$ACTIVE_SERVICE" ]; then
  FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr '
' ' ' || true)
  echo -e "\n"
  echo -e "DNS STATE: $FINAL_DNS"
  echo -e ""
fi

if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
  if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
    echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
    echo "  The gateway gracefully fell back to yesterday's cached blocklist."
    echo "  (Run 'docker logs rule-compiler' later to investigate which link died.)"
  fi
```

```
fi

echo -e "\n"
read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
  echo "Reversing gateway state..."
  exec bash "$0"
fi

echo -e "\nYou can close this terminal window now."

exit 0
fi

#
# macOS TOGGLE  jot / netstat / networksetup, launchd / pf.
#

# Enforce Cryptographic Script Integrity
if [ -f "scripts/crypto.sh" ]; then
  if ! bash scripts/crypto.sh --verify; then
    exit 1
  fi
fi

# Define log file
LOG_FILE="$PWD/output.log"

# Rotate log file if it exceeds 1GB (1073741824 bytes).
# Sized large on purpose: with DEBUG_TRACE the log holds a full command-by-command
# trace  effectively the entire compilation/warning history of the framework
# which is valuable to keep for post-mortems. Raised from 50MB so always-on
# tracing doesn't churn that history away prematurely.
if [ -f "$LOG_FILE" ]; then
  log_size=$(stat -f%z "$LOG_FILE" 2>/dev/null || stat -c%s "$LOG_FILE" 2>/dev/null || echo 0)
  if [ "$log_size" -gt 1073741824 ]; then
    mv "$LOG_FILE" "${LOG_FILE}.old" 2>/dev/null || rm -f "$LOG_FILE"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] Log file exceeded 1GB; rotated to output.log.old" > "$LOG_FILE"
  fi
fi

# --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
# TOGGLE_PHASE is updated as the script moves through STARTUP/STOP/START so that
# if the process is killed (window closed, pkill, OOM) or exits with an error,
# the EXIT trap records the final exit code AND the phase it died in. Previously
# an interrupted START left output.log ending abruptly after the teardown with no
# indication of what failed  this makes every exit traceable.
TOGGLE_PHASE="init"
_log_exit_breadcrumb() {
  local rc=$?
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
    2>/dev/null || true
}
trap '_log_exit_breadcrumb' EXIT

# --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
# When enabled, mirror a full command-by-command xtrace to output.log for a
# complete post-mortem trail (e.g. a race during --restart while Wi-Fi is
# re-associating). Off by default. Read directly from .env here because
# common.sh (read_env_var) isn't sourced yet.
#
# BASH_XTRACEFD (send xtrace to its own fd, keeping the terminal clean) needs
# bash >= 4.1. macOS ships /bin/bash 3.2, so we branch:
#   - bash >= 4.1 (e.g. Linux): dedicate fd 9 to the log and point xtrace there;
#     the terminal stays clean and normal stderr is untouched.
#   - bash 3.2 (stock macOS): BASH_XTRACEFD is unsupported, so xtrace always goes
#     to stderr. We redirect stderr *straight to the log file* with a plain
#     `exec 2>>"$LOG_FILE"`. This deliberately does NOT use a process
#     substitution (`2> >(tee )`): under `set -e`, stderr flowing through the
#     backgrounded `tee` proved to abort the script mid-START on 3.2 (a failed
#     write / SIGPIPE in a pipeline returned non-zero and tripped `set -e`, with
#     $LINENO mis-blaming the enclosing `fi`  see devref 15.11.8-A/E). The
#     tradeoff: on 3.2 the terminal no longer shows live progress while tracing
#     (everything is in output.log instead). Acceptable for an opt-in debug mode.
# We NEVER use the `exec {var}>>` dynamic-fd form (4.1+ only), which hard-fails on 3.2.
DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'" | tr '[:upper:]' '[:lower
    :]' | tr -d '[:space:]')
if [ "$DEBUG_TRACE" = "true" ]; then
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled  xtrace mirrored to output.log (bash ${BASH_VERSINFO
    [0]}.${BASH_VERSINFO[1]})" >> "$LOG_FILE"
  # Timestamped, location-aware xtrace prompt (works on 3.2 and 4.x).
```

```
  export PS4='+ [${SECONDS}s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: '  # subshell-free clock: a $(...) in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
  if [ "${BASH_VERSINFO[0]}" -gt 4 ] || { [ "${BASH_VERSINFO[0]}" -eq 4 ] && [ "${BASH_VERSINFO[1]}" -ge 1 ]; }; then
    exec 9>>"$LOG_FILE"
    export BASH_XTRACEFD=9
    set -x
  else
    # bash 3.2 (stock macOS): plain stderrfile redirect (no process substitution).
    exec 2>>"$LOG_FILE"
    set -x
  fi
fi

# --- MAIN EXECUTION LOGGING ---
echo -e "\n=======================================================" >> "$LOG_FILE"
echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh invoked" >> "$LOG_FILE"
echo "=======================================================" >> "$LOG_FILE"

# (Previously this script chmod 600'd .env at start to "unlock" it for docker compose,
# and chmod 000'd it on exit. That pattern broke docker compose on macOS/Colima when
# the umask produced 0600  the script removed its own ability to read .env before it could even proceed.)

# --- PROCEDURAL PORT GENERATION ---
# To prevent static port fingerprinting by malware, we randomize exposed ports
# on every boot and persist them to a hidden environment file.
PORTS_FILE="/tmp/nullexit-ports.env"
if [[ "$1" == "" || "$1" == "--restart" ]]; then

  # Collision-avoidance function
  GENERATED_PORTS=""
  get_free_port() {
    local port
    while true; do
      port=$(jot -r 1 10000 65000)
      # 1. Prevent self-collision within this loop
      if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
        continue
      fi
      # 2. Prevent collision with ANY process on the Mac (root daemons, terminal apps, etc)
      # We use netstat instead of lsof because netstat reads the kernel socket table
      # directly and doesn't require sudo to see ports bound by other users/root.
      if ! netstat -an -p tcp | awk '{print $4}' | grep -q "\.$port$"; then
        echo "$port"
        return
      fi
    done
  }

  export TOR_SOCKS_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_SOCKS_PORT"

  export TOR_CONTROL_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"

  export SOCKS_PROXY_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"

  export DNS_PROXY_PORT=$(get_free_port)

  echo "export TOR_SOCKS_PORT=$TOR_SOCKS_PORT" > "$PORTS_FILE"
  echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
  echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
  echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
  ADGUARD_PASS=$(grep -E "^ADGUARD_PASSWORD=" .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'")
  echo "export TOR_PASSWORD=${ADGUARD_PASS:-nullexit}" >> "$PORTS_FILE"
elif [ -f "$PORTS_FILE" ]; then
  # On STOP, source the existing ports so docker-compose down matches
  source "$PORTS_FILE"
else
  # Fallbacks if file is lost
  export TOR_SOCKS_PORT=9050
  export TOR_CONTROL_PORT=9051
  export SOCKS_PROXY_PORT=1080
  export DNS_PROXY_PORT=5354
fi

# Start execution timer
TOGGLE_START_TIME=$SECONDS

if [[ "$1" == "--restart" ]]; then
```

```
  echo "Executing Gateway Restart Sequence..."
  # We must source common.sh here temporarily to check gateway state before we proceed
  source "$SCRIPT_DIR/scripts/common.sh"
  # Use persisted intent (marker) via gateway_state, which echoes a string (never
  # a return code) to stay set -e-safe under set -x. See devref 15.11.8.
  if [ "$(gateway_state)" = "running" ]; then
    echo "Gateway is currently running. Stopping it first..."
    bash "$0"
    echo "Gateway stopped. Now starting it..."
  else
    echo "Gateway is already stopped. Starting it..."
  fi
  bash "$0"
  exit 0
fi

# Global flag to track if the script completed successfully
SUCCESS_RUN=false

# Global variable to track the currently running background command PID
CURRENT_BG_PID=""

# Source common bash functions
source "$SCRIPT_DIR/scripts/common.sh"

# Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
LOCK_FILE="/tmp/nullexit-toggle.lock"
if [ -f "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
      die "Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running."
    fi
  fi
fi
echo "$$" > "$LOCK_FILE"

# Helper function to run a GUI command as the logged-in console user
# (Prevents permission failures if the script is run with sudo)
run_gui_cmd() {
  local console_user
  console_user=$(stat -f '%Su' /dev/console 2>> output.log || echo "$SUDO_USER")
  if [ -z "$console_user" ] || [ "$console_user" = "root" ]; then
    console_user=$(logname 2>> output.log || echo "$USER")
  fi

  if [ -n "$console_user" ] && [ "$console_user" != "root" ] && [ "$EUID" -eq 0 ]; then
    sudo -u "$console_user" "$@"
  else
    "$@"
  fi
}

# Active-state marker: written at end of START path so external watchers
# (see launchd/com.nullexit.wake-recovery.plist + scripts/watcher.sh +
# devref 10.29) know the gateway is currently up. They only fire
# recover.sh --post-wake when this file is present. Cleared at the top
# of the STOP path and inside cleanup_handler on any error/signal path
# so a half-broken gateway never re-triggers post-wake refreshes.
write_gateway_active_marker() {
  date -u +%FT%TZ > /tmp/nullexit-gateway-active.marker
  # Set the watcher debounce timestamp to now to prevent a redundant post-startup recovery run.
  date +%s > /tmp/nullexit-watcher.last-recovery 2>/dev/null || true
}

clear_gateway_active_marker() {
  rm -f /tmp/nullexit-gateway-active.marker
  rm -f "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/TUNNEL_FAILED_CLOSED.marker"
}

# Capture whether the gateway-active marker exists BEFORE the defensive clear
# below wipes it. The STOP-vs-START decision (gateway_should_be_running) reads
# this captured value  persisted *intent*  instead of a live Docker probe, so
# a disable is instant and works even when Colima/Docker is down. Must be read
# before clear_gateway_active_marker or it would always look "stopped".
#
# Written as default-assign + `if` WITHOUT an `else`: under `set -e` (with the
# DEBUG_TRACE xtrace redirect active) an `if [ -f  ]; then ; else ; fi` whose
# condition is false was aborting the script here at init on macOS /bin/bash 3.2
#  the false `[ -f ]` status leaked through the compound. An `if` with no `else`
# yields status 0 when the condition is false, so nothing leaks.
```

```
GATEWAY_MARKER_AT_STARTUP="no"
if [ -f "$MARKER_FILE" ]; then
  GATEWAY_MARKER_AT_STARTUP="yes"
fi

# Defensive: clear any stale marker from a prior crashed/aborted run. If a
# previous toggle.sh wrote the marker but was OOM-killed in the middle of the
# START block (before the END-of-START write)  or if the user rebooted mid
# session  this prevents the watcher from firing post-wake against a
# non-running gateway on every wake event. Placed AFTER the function
# definitions so bash resolves the call at parse time.
clear_gateway_active_marker

PID_FILE="$PID_CAFFEINATE"

# Start macOS caffeinate to prevent sleep while gateway is running
start_sleep_prevention() {
  if ! command -v caffeinate >/dev/null 2>&1; then
    echo "  [Warning] caffeinate tool not found. Sleep prevention unavailable."
    return 1
  fi

  # Stop any existing sleep prevention process first to avoid duplicates
  stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"

  echo "  Enabling system sleep prevention (caffeinate)..."
  # Run caffeinate wrapped in a bash trap to prevent idle system sleep.
  # Critically, this traps macOS shutdown signals (SIGTERM) and automatically
  # flushes hijacked DNS back to normal right before the Mac powers off.
  # We do not use recover.sh here because it is a heavy recovery script (managing
  # containers, restarting tailscaled, etc.) which takes too long and gets
  # terminated by macOS before it can reset the DNS. Instead, we surgically and
  # instantly restore normal DNS and proxy settings here.
  #
  # -i alone only blocks *idle* sleep  a stopped/crashed gateway process,
  # a hung docker command, or any other assertion lapse still lets the Mac
  # sleep and drop every mesh device's exit node. -s additionally blocks
  # sleep outright whenever the Mac is on AC power (it's a no-op on battery,
  # so unplugged runs still sleep normally and don't burn charge).
  nohup bash -c "
    trap '
      echo \"[Shutdown Trap] Reverting network settings...\" >> \"$PWD/output.log\" 2>&1
      kill \$! 2>/dev/null || true
      sudo -n networksetup -setdnsservers \"$ACTIVE_SERVICE\" 1.1.1.1 >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setdnsservers \"$EN0_SERVICE\" 1.1.1.1 >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsearchdomains \"$ACTIVE_SERVICE\" \"Empty\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsearchdomains \"$EN0_SERVICE\" \"Empty\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsocksfirewallproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 ||
     true
      sudo -n networksetup -setsocksfirewallproxystate \"$EN0_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setwebproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setwebproxystate \"$EN0_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsecurewebproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsecurewebproxystate \"$EN0_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      if command -v tailscale >/dev/null 2>&1; then
        tailscale up >> \"$PWD/output.log\" 2>&1 || true
        tailscale set --accept-dns=false --exit-node= >> \"$PWD/output.log\" 2>&1 || true
        TS_DOWN_ARGS=\"\"
        if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
        tailscale down \$TS_DOWN_ARGS >> \"$PWD/output.log\" 2>&1 &
      fi
      sleep 1
      echo \"[Shutdown Trap] Cleanup complete.\" >> \"$PWD/output.log\" 2>&1
      exit 0
    ' SIGTERM SIGINT SIGHUP
    caffeinate -i -s &
    wait \$!
  " >> output.log 2>&1 &
  local caffe_pid=$!
  echo "$caffe_pid" > "$PID_FILE"
  echo "  Sleep prevention active (PID $caffe_pid). Your Mac won't sleep while the gateway is running."
}



DNS_WATCHER_PID_FILE="$PID_DNS_WATCHER"
WARP_WATCHER_PID_FILE="/tmp/nullexit-warp-watcher.pid"



# Background WARP tunnel liveness monitor. Polls cdn-cgi/trace every 30s
```

```
# while healthy, but accelerates to polling every 5s if a failure is detected.
# Always logs state transitions to output.log. When warp=off persists for
# WARP_FAIL_THRESHOLD consecutive polls (default 6 = 30s of downtime), the watcher
# triggers a nuclear recovery: runs recover.sh to tear down the entire
# gateway  disconnecting Tailscale, stopping containers, resetting DNS,
# and power-cycling Wi-Fi. Note: This minimizes exposure window, but
# only the pf kill switch guarantees absolutely zero leakage on drop.
#
# The threshold (default 6) is statistically chosen: even at an
# unrealistically high 10% false-positive rate per poll, P(6 consecutive
# false positives) = 0.1^6  0.0001%. Real-world false-positive rates
# are far lower, making 30s of continuous downtime overwhelmingly likely
# to be a genuine outage.
#
# Configurable via .env:  WARP_FAIL_THRESHOLD=3  (15s, more aggressive)
start_warp_watcher() {
  stop_warp_watcher
  # Parse threshold from .env (same pattern as GATEWAY_MSS, GATEWAY_BYPASS_PING, etc.)
  local threshold
  threshold=$(read_env_var WARP_FAIL_THRESHOLD)
  if [ -z "$threshold" ]; then
    threshold="${WARP_FAIL_THRESHOLD:-6}"
  fi
  local trigger_file="/tmp/nullexit-warp-shutdown-triggered"
  rm -f "$trigger_file"
  # Resolve the physical uplink interface (Wi-Fi/Ethernet) once, same convention
  # as wait_for_dhcp_settle(). The watcher uses it to tell "no uplink at all"
  # (e.g. Wi-Fi dropped with auto-join off) apart from a genuine WARP failure,
  # so a transient link loss can't false-trigger the nuclear gateway teardown.
  local phys_iface
  phys_iface=$(networksetup -listallhardwareports 2>/dev/null | awk '/Hardware Port: (Wi-Fi|Ethernet)/{getline;
      print $2; exit}')
  [ -z "$phys_iface" ] && phys_iface="en0"
  echo "  Starting background WARP Watcher (polling every 30s [accelerating to 5s on failure], shutdown after $
     {threshold} consecutive failures; paused while ${phys_iface} has no uplink)..."
  nohup bash -c "
    trap 'exit 0' SIGTERM SIGINT SIGHUP
    last_state='on'
    consec_off=0
    paused=0
    threshold='$threshold'
    trigger_file='$trigger_file'
    phys_iface='$phys_iface'
    out_log='$PWD/output.log'
    recover_bin='$PWD/recover.sh'
    inhibit_file='/tmp/nullexit-warp-inhibit.marker'
    wifi_helper='$PWD/scripts/wifi-rejoin.sh'
    while true; do
      #  Inhibit check
      # recover.sh --post-wake writes this marker while force-recreating the
      # warp container. During that window, the container is intentionally
      # down  don't count it as a failure or we'll fire nuclear recover.sh
      # and kill a gateway that was in the process of healing itself.
      if [ -f \"\$inhibit_file\" ]; then
        consec_off=0
        sleep 5
        continue
      fi

      #  Physical-uplink + link-state check
      # If the Mac has no *working* uplink (Wi-Fi dropped, or an 802.1x
      # de-auth/roam left en0 holding a stale DHCP lease it can no longer
      # route through), the cdn-cgi/trace probe below reads warp=off only
      # because there's no internet  NOT because the WARP tunnel failed.
      # The PF kill-switch means there's no egress path to leak through
      # either, so PAUSE failure-counting until the link returns instead of
      # nuking a gateway that heals itself once Wi-Fi reconnects. (A working
      # link + warp=off is a genuine failure and still trips the shutdown
      # below, unchanged.)
      #
      # 'Has an IP' is NOT 'has egress': on an enterprise-Wi-Fi de-auth macOS
      # keeps the stale lease on en0 for a while  so a bare getifaddr check
      # passed while the link was already dead and the gateway got nuked
      # anyway (recurred 2026-07-23; see devref 15.6.2 pt 5 / 15.12.27).
      # After confirming a routable IP we therefore also read the driver's
      # link-layer association state (\`ifconfig ... status:\`). We deliberately
      # do NOT ping the router: on this AP-isolated, kill-switched network the
      # gateway drops ICMP even when the tunnel is perfectly healthy, so a
      # ping probe would false-PAUSE a live gateway. 'status: active' =
      # associated; 'inactive' = link down under a stale lease. Fail OPEN
      # (trust the lease) for any other/absent status so an unexpected driver
```

```
      # string can't strand the watcher.
      phys_ip=\$(ipconfig getifaddr \"\$phys_iface\" 2>/dev/null)
      uplink_dead=0
      case \"\$phys_ip\" in
        ''|169.254.*)
          uplink_dead=1 ;;                        # no routable IP at all
        *)
          link_status=\$(ifconfig \"\$phys_iface\" 2>/dev/null | awk '/status:/{print \$2; exit}')
          case \"\$link_status\" in
            inactive|none) uplink_dead=1 ;;        # stale lease over a dead link
          esac ;;
      esac
      if [ \"\$uplink_dead\" = 1 ]; then
        if [ \"\$paused\" = 0 ]; then
          echo \"[\$(date -u +%FT%TZ)] WARP watcher PAUSED  no working uplink on \$phys_iface (link down / not associated, or stale lease). Not counting warp=off as a failure; gateway left intact until the link returns.\" >> \"\$out_log\"
          paused=1
        fi
        # Re-associate ONLY the last-known-good network (never any other),
        # after a grace period so a DHCP-only blip can self-heal first.
        bash \"\$wifi_helper\" rejoin >/dev/null 2>&1 || true
        consec_off=0
        sleep 10
        continue
      fi
      if [ \"\$paused\" = 1 ]; then
        echo \"[\$(date -u +%FT%TZ)] WARP watcher RESUMED  working uplink back on \$phys_iface (\$phys_ip). Normal WARP monitoring restored.\" >> \"\$out_log\"
        paused=0
        # Link is back  clear the drop-timer so the next drop re-arms the grace window.
        bash \"\$wifi_helper\" reset >/dev/null 2>&1 || true
      fi

      state='off'
      if docker compose exec -T warp wget -qO- --timeout=5 \
           https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
           | grep -q 'warp=on'; then
        state='on'
      fi

      #  State-transition logging
      if [ \"\$state\" != \"\$last_state\" ]; then
        if [ \"\$state\" = 'off' ]; then
          echo \"[\$(date -u +%FT%TZ)] WARP DOWN  cdn-cgi/trace reports warp=off\" >> \"\$out_log\"
        fi
        last_state=\"\$state\"
      fi

      #  Consecutive-failure tracking + auto-shutdown
      if [ \"\$state\" = 'off' ]; then
        consec_off=\$((consec_off + 1))
        if [ \"\$consec_off\" -ge \"\$threshold\" ] && [ ! -f \"\$trigger_file\" ]; then
          touch \"\$trigger_file\"
          echo \"[\$(date -u +%FT%TZ)] WARP SHUTDOWN  \$consec_off consecutive failures (threshold=\$threshold). Running recover.sh to kill the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.\" >> \"\$out_log\"
          # Nuclear recovery: tear down the entire gateway.
          # recover.sh resets DNS  1.1.1.1, disconnects Tailscale,
          # stops Docker containers, flushes routes, power-cycles Wi-Fi.
          bash \"\$recover_bin\" --auto >> \"\$out_log\" 2>&1 || true
          echo \"[\$(date -u +%FT%TZ)] WARP SHUTDOWN  recover.sh completed, watcher exiting. Your IP is now your real ISP IP.\" >> \"\$out_log\"
          exit 0
        fi
      else
        if [ \"\$consec_off\" -gt 0 ]; then
          echo \"[\$(date -u +%FT%TZ)] WARP RECOVERED  warp=on after \$consec_off consecutive off readings (threshold was \$threshold)\" >> \"\$out_log\"
        fi
        consec_off=0
      fi

      if [ \"\$state\" = \"off\" ]; then
        sleep 5
      else
        sleep 30
      fi
    done
  " >> output.log 2>&1 &
  echo $! > "$WARP_WATCHER_PID_FILE"
```

```
}

stop_warp_watcher() {
  if [ -f "$WARP_WATCHER_PID_FILE" ]; then
    local wp
    wp=$(cat "$WARP_WATCHER_PID_FILE")
    if [ -n "$wp" ] && kill -0 "$wp" 2>/dev/null; then
      echo "  Stopping background WARP Watcher (PID $wp)..."
      kill "$wp" 2>/dev/null || true
    fi
    rm -f "$WARP_WATCHER_PID_FILE"
  fi
}




# Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
cleanup_handler() {
  local exit_code=$?
  local trigger_type="$1"
  local line_no="$2"

  if [ "$SUCCESS_RUN" = "false" ]; then
    echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1...
    "
    if [ -n "$ACTIVE_SERVICE" ]; then
      reset_dns
    fi

    if [ -n "$CURRENT_BG_PID" ]; then
      echo "Terminating active command (PID $CURRENT_BG_PID)..."
      kill -15 "$CURRENT_BG_PID" 2>> output.log || true
    fi

    # Disconnect host Tailscale and close the GUI application on failure
    disconnect_tailscale_host

    # Stop local DNS proxy if running
    stop_local_dns_proxy

    # Stop sleep prevention
    stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
    stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
    stop_warp_watcher
    clear_gateway_active_marker

    # Capture warp logs on failure for debugging before teardown
    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n--- WARP FAILURE LOGS ---" >> output.log
      docker compose logs warp --tail=100 >> output.log 2>&1 || true
      echo "-------------------------" >> output.log
    fi

    # Best-effort container teardown on failure
    docker compose down --remove-orphans -t 30 2>> output.log || true

    # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
    if [ -n "$ACTIVE_SERVICE" ]; then
      cleanup_network_state
    fi


    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n============================================="
      echo "ERROR: Script failed at line $line_no with exit code $exit_code."
      echo "============================================="
      echo "Troubleshooting steps:"
      echo "1. Run 'colima status' to verify the VM is active."
      echo "2. Run 'docker compose ps' to check container health."
      echo "3. Run 'docker compose logs' to view application logs."
      echo "4. Run 'tailscale status' to check host Tailscale status."
      echo "============================================="
    fi
    rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
  fi
}

trap 'cleanup_handler ERR $LINENO' ERR
trap 'cleanup_handler INT' INT
```

```
trap 'cleanup_handler TERM' TERM
trap 'cleanup_handler HUP' HUP

#  NOPASSWD sudo
# The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
# needed (networksetup, pfctl, route, ifconfig, dscacheutil, killall, pkill,
# and python3 -I scripts/dns-proxy.py for the fallback DNS proxy).
# All privileged calls below use `sudo -n` which will run without prompting.
# No credential caching loop is needed.

# Add Homebrew and standard paths
setup_standard_path

# Check for local DNS proxy tools (used when tailnet data plane is unavailable)
# Uses a Python DNS proxy (handles TCP wire format correctly).
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_BIN=""
if command -v python3 >> output.log 2>&1; then
  DNS_PROXY_BIN="python3 -I"
elif command -v python >> output.log 2>&1; then
  DNS_PROXY_BIN="python -I"
fi

# Path to the DNS proxy script (sibling of this script)
DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/dns-proxy.py"

# Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
# (`brew install tailscale` + `brew services start tailscale`)  NOT the
# Tailscale.app GUI  so there is no .app bundle or Network Extension sandbox
# to launch, no menu-bar click to perform, and no scutil Network Service to
# start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
# if you ran the install with sudo) and the control socket is always available.
TS_BIN=""
if command -v tailscale >> output.log 2>&1; then
  TS_BIN="tailscale"
fi

# Baseline: disable Tailscale DNS management at the daemon level.
# `tailscale set` writes a persistent preference. However, `tailscale up --reset`
# (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
# re-enable `--accept-dns=true` and nullify this `set`.
# Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
# `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
#
# This initial `set` is still useful as a baseline for non-reset operations
# (e.g. if the user runs `tailscale up` manually without --reset) and covers
# the brief window between this line and the first `--reset` call.
#
# Runs once at script start while tailscaled is still connected (if it was
# left running from a previous toggle), so the pref takes effect before any
# disconnection or reconnection logic.
if [ -n "$TS_BIN" ]; then
  $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
fi

# Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
# *entire* teardown  no scutil tunnel stop, no Tailscale.app GUI quit, no pkill.
# tailscaled keeps running (as a LaunchAgent / LaunchDaemon), which is intentional:
# the next `tailscale up` then reconnects in seconds rather than waiting for a
# slow .app launch + Network Extension activation. The `tailscale down` call also
# retains the user's auth state, so we don't force a browser re-login every cycle.
#
# Defined early (before the if/else branches and referenced by cleanup_handler's
# ERR trap) so that any failure before the main logic can still tear down Tailscale.


#  Wait for Network Readiness
# On --restart, the gateway tears down while the host network interface may
# still be reconnecting (Wi-Fi, Ethernet, USB tethering, etc.). If we proceed
# too early, get_active_service returns nothing, routing targets the wrong
# interface, and DNS/WARP setup silently fails.
#
# We use `route get default` to detect a valid default gateway rather than
# checking a specific interface (e.g. en0)  this generalizes across all
# connection types: Wi-Fi, Ethernet, USB-C adapters, iPhone tethering, etc.
# The moment the OS has any routable default gateway, we proceed.
_net_wait_secs=0
_net_timeout=60
echo "Waiting for network connectivity before starting..."
while true; do
  if route get default 2>/dev/null | grep -q 'gateway:'; then
```

```
    _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
    echo "  Network ready (default gateway: $_gw)."
    break
  fi
  if [ "$_net_wait_secs" -ge "$_net_timeout" ]; then
    echo "  [Warning] No default gateway after ${_net_timeout}s  proceeding anyway."
    break
  fi
  sleep 2
  _net_wait_secs=$((_net_wait_secs + 2))
done

ACTIVE_SERVICE=$(get_active_service)

# Resolve the service name for en0 (usually "Wi-Fi")  macOS scutil DNS resolver
# is commonly scoped to en0, so per-service DNS changes on other interfaces are
# ignored unless en0's service is also updated.
EN0_SERVICE=$(get_en0_service)

# Helper to add host-side bypass routes for the WARP WireGuard endpoints.
# This prevents an infinite recursive routing loop (tunnel loop) where
# the container's WireGuard packets exit the VM, reach the host, match the
# default route (the exit node), and get routed back into the tunnel.
# We route via the gateway IP (not interface) to ensure packets are routed
# properly even when the host's default route changes to the Tailscale interface.

# Helper to manually override the host's default route to point through the
# Tailscale utun* interface. Standalone tailscaled on macOS (Homebrew version)
# lacks the platform-specific integration to override the default gateway
# automatically when --exit-node is used.
setup_exit_node_routing() {
  local ts_iface
  ts_iface=$(ifconfig | awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':')

  if [ -n "$ts_iface" ]; then
    echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
    local host_mtu
    host_mtu=$(read_env_var HOST_MTU)
    host_mtu=${host_mtu:-1200}

    # Lower the host Tailscale MTU (defaults to 1200) to prevent fragmentation when passing
    # through the container's WARP tunnel (which also has an MTU of 1280).
    sudo -n ifconfig "$ts_iface" mtu "$host_mtu" >> output.log 2>&1 || true

    # DO NOT delete the physical default route! It crashes Colima.
    # Use the /1 trick to mathematically override the default route.
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true

    if netstat -nr | grep -E -q "^(0\.0\.0\.0/1|0/1)[[:space:]]+.*$ts_iface"; then
      echo "  [] Default route successfully overridden via $ts_iface."
    else
      echo "  [!] Failed to verify exit node routing on $ts_iface."
    fi
  else
    echo "[Warning] Could not detect Tailscale utun interface for host routing."
  fi
}

# Helper to nuke leftover network state after teardown (proxy settings, routing
# table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
# problem where stale state kills internet after the gateway stops.
cleanup_network_state() {
  echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."

  # 1. Disable any leftover proxy settings (needs root on many macOS versions)
  for svc in "$ACTIVE_SERVICE" "$EN0_SERVICE"; do
    disable_all_proxies "$svc"
  done
  echo "  Proxies disabled."

  # 2. Flush DNS cache
  if command -v dscacheutil >> output.log 2>&1; then
    sudo -n dscacheutil -flushcache 2>> output.log || true
  fi
  sudo -n killall -HUP mDNSResponder 2>> output.log || true
  echo "  DNS cache flushed."

  # 3. Flush stale routing table entries
```

```
  if command -v route >> output.log 2>&1; then
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
  fi
  remove_warp_bypass_routes
  remove_apple_split_routes
  disable_killswitch
  echo "  Routing table and firewall flushed."

  # 4. Power-cycle Wi-Fi to clear any lingering interface state
  WIFI_PORT=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print $2}')
  if [ -n "$WIFI_PORT" ]; then
    sudo -n networksetup -setairportpower "$WIFI_PORT" off 2>> output.log || true
    sleep 2
    sudo -n networksetup -setairportpower "$WIFI_PORT" on 2>> output.log || true
    echo "  Wi-Fi power-cycled (interface $WIFI_PORT)."
    sleep 3
  else
    echo "  Wi-Fi interface not detected; bouncing en0..."
    sudo -n ifconfig en0 down 2>> output.log || true
    sudo -n ifconfig en0 up 2>> output.log || true
  fi

  # Force restart sharing services to prevent AirDrop / AirPlay discovery freezes
  # after network interface/DNS changes.
  echo "  Resetting macOS sharing services (AirDrop/AirPlay)..."
  reset_sharing_services

  echo "Network state cleanup complete."
}

SOCKS_PROXY_PORT=${SOCKS_PROXY_PORT}

# Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
# during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
# `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
# would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.


echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
reset_dns




echo "Checking Gateway Status..."
HIJACK_HOST=$(read_env_var GATEWAY_HIJACK_HOST | tr '[:upper:]' '[:lower:]')

# Decide STOP vs START from persisted intent (marker), not a live Docker probe.
# This makes *disable* instant and robust even when Colima/Docker is down.
#
# CRITICAL (set -e + set -x on bash 3.2): gateway_state ECHOES "running"/"stopped"
# and always returns 0, so we branch on the STRING  never on a function's exit
# code. A helper that `return 1`s trips a spurious `set -e` abort under `set -x`
# (DEBUG_TRACE=true) even when guarded, hence the echo contract. See devref 15.11.8.
if [ "$(gateway_state)" = "running" ]; then
  TOGGLE_PHASE="stop"
  echo -e "\n=============================================="
  echo "Gateway is RUNNING. Stopping it now..."
  echo -e "==============================================\n"

  # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
  if [[ "$HIJACK_HOST" != "false" ]]; then
    disconnect_tailscale_host
    echo ""
  fi

  echo "Stopping Docker containers..."
  # `|| true`: a disable must always complete its teardown even if the Docker
  # daemon / Colima VM is already down (compose down would exit non-zero and,
  # under set -e, abort the STOP). The START path intentionally does NOT guard
  # its `docker compose up`  a failed start SHOULD trigger cleanup.
  run_logged docker compose down --remove-orphans -t 30 || true

  # Only stop Colima if the user explicitly opted in (false by default) to prevent breaking their other Docker dev projects
  STOP_COLIMA=$(read_env_var STOP_COLIMA_ON_EXIT | tr '[:upper:]' '[:lower:]')
  if [[ "$STOP_COLIMA" == "true" ]]; then
    echo -e "\nStopping Colima VM to free up host RAM and battery..."
    run_with_timeout 30 colima stop >> output.log 2>&1 || echo "Warning: Failed to stop Colima gracefully."
```

```
  else
    echo -e "\nLeaving Colima running. (Set STOP_COLIMA_ON_EXIT=true in .env to change this)"
  fi

  if [[ "$HIJACK_HOST" != "false" ]]; then
    # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
    # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
    # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
    echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
    reset_dns

    # 3. Stop local DNS proxy if running
    stop_local_dns_proxy

    # Stop background daemons
    stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
    stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
    stop_warp_watcher
    clear_gateway_active_marker

    # 4. Nuke leftover network state so internet actually works after teardown
    cleanup_network_state
  fi

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
else
  TOGGLE_PHASE="start"
  START_GATEWAY=true
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] Action: STARTUP GATEWAY" >> "$LOG_FILE"
  echo -e "\n=============================================="
  echo "Gateway is STOPPED. Starting it now..."
  echo -e "==============================================\n"

  # 1. Prevent host exit-node deadlock during VM / Container startup
  if [[ "$HIJACK_HOST" != "false" ]]; then
    disconnect_tailscale_host
  fi

  # 2. Wait for physical DHCP lease to settle (crucial for restarts/wake-up/roaming)
  wait_for_dhcp_settle

  # 3. Boot Colima VM if it is not already running
  echo -e "\nChecking Colima VM status..."
  if ! run_with_timeout 15 colima status >> output.log 2>&1; then
    # Colima networking is gated on the SAME LAN-P2P signal routing-fix uses
    # (.lan_p2p_detected overrides TAILSCALE_ALLOW_LAN_P2P in .env). A LAN-reachable
    # VM (--network-address + bridged/shared) needs the socket_vmnet helper and is
    # only worthwhile on a trusted home network, so we request it only when P2P is
    # allowed; otherwise  AP-isolated / enterprise, the common case  we boot plain
    # user-mode networking. Either way host<->container P2P is UNAFFECTED: it rides
    # routing-fix's Docker-gateway / .host_ips RETURN rules, not the VM network mode
    # (see devref 15.3.5). colima_start_until_ready returns as soon as Docker answers
    # (~15s) instead of blocking on colima's ~2min post-provision hang (see 15.8.7);
    # that hang is colima-internal and mode-independent (it also occurs in user mode).
    allow_p2p=$(read_env_var "TAILSCALE_ALLOW_LAN_P2P")
    [ -z "$allow_p2p" ] && allow_p2p="false"
    if [ -f "$SCRIPT_DIR/.lan_p2p_detected" ]; then
      allow_p2p=$(tr -d '[:space:]' < "$SCRIPT_DIR/.lan_p2p_detected" 2>/dev/null || echo "false")
    fi

    if [ "$allow_p2p" = "true" ]; then
      colima_network_mode="bridged"
      if [[ "$OSTYPE" == "darwin"* ]]; then
        security_type=$(system_profiler SPAirPortDataType 2>/dev/null | awk '/Current Network Information:/{
    found=1} found && /Security:/{print $NF; exit}')
        if echo "$security_type" | grep -qi "Enterprise"; then
          echo -e "\n  [!] WPA2-Enterprise Wi-Fi detected! Bridged mode will cause MAC-spoofing
    deauthentication."
          echo "      Falling back to '--network-mode shared' for Colima."
          colima_network_mode="shared"
        fi
      fi
      echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P on: --network-address
    $colima_network_mode)..."
      colima_start_until_ready --memory 0.6 --vm-type vz --network-address --network-mode "$colima_network_mode
    " \
        || echo "  [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry
    )."
    else
      colima_network_mode="user"
```

```
      echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P off: fast user-mode networking) ..."
      colima_start_until_ready --memory 0.6 --vm-type vz \
        || echo "  [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry)."
    fi
    echo "$colima_network_mode" > /tmp/nullexit_colima_mode.txt
  else
    echo "Colima is already running."
  fi

  # 3b. Auto-detect and fix Docker Subnet Collisions (e.g. if the user travels to a 172.17.x.x Wi-Fi)
  if [ -f "scripts/fix-docker-bridge-collision.sh" ]; then
    if bash scripts/fix-docker-bridge-collision.sh --dry | grep -q "collision detected"; then
      echo -e "\n[!] WARNING: Docker Subnet Collision Detected with your current Wi-Fi!"
      echo -e "    Auto-fixing Colima config to prevent internet jitter..."
      bash scripts/fix-docker-bridge-collision.sh --skip-toggle || true
    fi
  fi

  # 3c. Configure swap file inside the VM to prevent OOM on low-memory limits
  # We also set vm.swappiness=10 so the Linux kernel strictly prefers physical RAM
  # and avoids unnecessarily wearing out the SSD with proactive background swapping.
  if ! run_with_timeout 15 colima ssh -- grep -q 'swapfile' /proc/swaps >> output.log 2>&1; then
    echo "Configuring 400MB swap file inside the VM to prevent OOM..."
    run_with_timeout 30 colima ssh -- sudo sh -c "if [ ! -f /swapfile ]; then dd if=/dev/zero of=/swapfile bs=1M count=400 status=none && mkswap /swapfile; fi && swapon /swapfile && sysctl vm.swappiness=10" >> output.log 2>&1 || echo "Warning: Failed to enable swap file inside the VM."
  fi

  # 4. Clean up corrupted AdGuardHome configurations
  if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
    rm -f "adguard/conf/AdGuardHome.yaml"
    echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
  fi

  # 4b. Auto-heal stale Docker socket from hard reboots
  if ! run_with_timeout 15 docker ps >/dev/null 2>&1; then
    echo "Docker daemon is unresponsive (likely a stale socket from a crash). Auto-healing Colima..."
    run_with_timeout 120 colima restart >> output.log 2>&1
  fi

  # 4c. Inject TCP MSS clamp from .env into post-rules.txt before starting
  if [ -f .env ] && grep -q "^GATEWAY_MSS=" .env; then
    GATEWAY_MSS=$(read_env_var GATEWAY_MSS)
    if [[ "$GATEWAY_MSS" =~ ^[0-9]+$ ]]; then
      # macOS sed syntax
      sed -i '' "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null || \
      # Linux sed fallback
      sed -i "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null
    fi
  fi

  # 5. Start compose services
  write_host_ips

  # Procedurally generated ports have already been exported at the top of the script.
  echo "  [Security] Procedurally randomized proxy ports generated to evade fingerprinting."

  # --- HONEY-PORT TRIPWIRE ---
  # Generate a random trap port. If any malware does a sequential port scan on localhost,
  # it will inevitably hit this port. When it does, we instantly fire a macOS alert.
  # Reap the honey-port from any previous toggle before arming a new one, so the
  # `nc -l` tripwire never accumulates one idle listener per toggle-on (15.12.20).
  pkill -f "nc -l localhost" 2>/dev/null || true
  HONEY_PORT=$(get_free_port)
  (
    if nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$LOG_FILE"
      osascript -e 'display notification "An unknown application is aggressively scanning your local ports. Malware port-scan suspected." with title "nullexit OPSEC Alert" sound name "Basso"' 2>/dev/null || true
    fi
  ) &
  disown %1 2>/dev/null || true
  echo "  [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."

  echo -e "\nStarting Docker containers..."

  # --- Ad-block rule compilation: hash-gated one-off, off the critical path (15.8.8) ---
  # The rule-compiler (a `compile`-profiled container, the SOLE writer of adguard/work)
```

```
  # re-dedupes ~340k rules inside the 600MB VM (~28s)  far too slow to run every toggle.
  # It is gated on a hash of the lists YOU control (black_list.txt + white_list.txt):
  # recompile only when they change, when compiled_rules.txt / ip_blocklist.ipset are
  # missing, or when they exceed RULES_MAX_AGE_DAYS (default 7, so the remote blocklists
  # still refresh occasionally). Otherwise AdGuard/routing-fix boot on the existing
  # artifacts and we skip the ~28s  they no longer `depends_on` the rule-compiler
  # container (see docker-compose.yml). A compile failure is non-fatal: we keep the
  # previous compiled rules rather than bricking the tunnel (ad-block != the kill-switch).
  RULES_HASH_FILE="$SCRIPT_DIR/.rules_hash"
  rules_hash=$(compute_rules_hash)
  stored_rules_hash=""
  [ -f "$RULES_HASH_FILE" ] && stored_rules_hash=$(cat "$RULES_HASH_FILE" 2>/dev/null)
  rules_max_age_days=$(read_env_var RULES_MAX_AGE_DAYS)
  [ -z "$rules_max_age_days" ] && rules_max_age_days=7
  need_compile=false
  if [ ! -f adguard/work/userfilters/compiled_rules.txt ] || [ ! -f adguard/work/userfilters/ip_blocklist.ipset ]; then
    need_compile=true
  elif [ -z "$rules_hash" ] || [ "$rules_hash" != "$stored_rules_hash" ]; then
    need_compile=true
  elif [ -n "$(find adguard/work/userfilters/compiled_rules.txt -mtime +"$rules_max_age_days" 2>/dev/null)" ]; then
    need_compile=true
  fi
  if [ "$need_compile" = "true" ]; then
    echo "  Ad-block lists changed (or artifacts missing/stale)  compiling rules (one-off, ~28s)..."
    if run_logged docker compose run --build --rm rule-compiler >> output.log 2>&1; then
      if [ -n "$rules_hash" ]; then echo "$rules_hash" > "$RULES_HASH_FILE"; fi
    else
      echo "  [!] Rule compile failed  continuing with existing compiled ad-block rules."
    fi
  else
    echo "  Ad-block lists unchanged  reusing compiled rules (skipping the ~28s recompile)."
  fi

  # --- Gate `--build` for the long-running images (routing-fix, tor) ---
  # BuildKit re-checking these on the 600MB VM costs ~30s even when every layer is
  # CACHED (15.8.8). Only pass --build when the hashed inputs change or an image is
  # missing; otherwise plain `up -d` reuses cached images. `up` starts everything
  # except the `compile`-profiled rule-compiler.
  BUILD_HASH_FILE="$SCRIPT_DIR/.build_hash"
  build_inputs_hash=$(compute_build_inputs_hash)
  stored_build_hash=""
  [ -f "$BUILD_HASH_FILE" ] && stored_build_hash=$(cat "$BUILD_HASH_FILE" 2>/dev/null)
  local_images_present=true
  for _img in nullexit-routing-fix:v1.0.0 nullexit-tor:v1.0.0; do
    docker image inspect "$_img" >> output.log 2>&1 || local_images_present=false
  done
  if [ -n "$build_inputs_hash" ] && [ "$build_inputs_hash" = "$stored_build_hash" ] && [ "$local_images_present" = "true" ]; then
    echo "  Build inputs unchanged  skipping image build (reusing cached images)."
    run_logged docker compose up -d
  else
    echo "  Build inputs changed or a local image is missing  building images..."
    run_logged docker compose up -d --build
    # Record the hash only after a successful build (set -e aborts above on failure).
    if [ -n "$build_inputs_hash" ]; then echo "$build_inputs_hash" > "$BUILD_HASH_FILE"; fi
  fi

  RULE_COUNT=$(grep "! Total Block Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -F': ' '{print $2}' || echo "0")
  NATIVE_COUNT=$(grep "! Native AdGuard Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -F': ' '{print $2}' || echo "0")

  if [ "$RULE_COUNT" != "0" ] && [ "$RULE_COUNT" != "" ]; then
    if [ "$NATIVE_COUNT" != "0" ] && [ "$NATIVE_COUNT" != "" ]; then
      TOTAL_COUNT=$((RULE_COUNT + NATIVE_COUNT))
      # Format with commas for readability (e.g. 441,578)
      if command -v printf &> /dev/null; then
        FORMATTED_RULE=$(printf "%'d" "$RULE_COUNT")
        FORMATTED_TOTAL=$(printf "%'d" "$TOTAL_COUNT")
        echo "  DNS: Compiled $FORMATTED_RULE unique custom rules (Total active in AdGuard: ~$FORMATTED_TOTAL rules)."
      else
        echo "  DNS: Compiled $RULE_COUNT unique custom rules (Total active in AdGuard: ~$TOTAL_COUNT rules)."
      fi
    else
      echo "  DNS: Compiled and loaded $RULE_COUNT optimized rules."
    fi
  fi

```

```
  # Report IP blocklist compilation results
  IP_COUNT=$(grep "^# Entries:" adguard/work/userfilters/ip_blocklist.ipset 2>/dev/null | awk -F': ' '{print $2
    }' || echo "0")
  if [ "$IP_COUNT" != "0" ] && [ "$IP_COUNT" != "" ]; then
    if command -v printf &> /dev/null; then
      FORMATTED_IP=$(printf "%'d" "$IP_COUNT")
      echo "  IP:  Loaded $FORMATTED_IP threat intelligence IPs/CIDRs into kernel firewall."
    else
      echo "  IP:  Loaded $IP_COUNT threat intelligence IPs/CIDRs into kernel firewall."
    fi
  fi

  # 6. Wait for the gateway container's Tailscale connection to be ready
  BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
  USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')

  echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
  echo "  (This can take 30-60s  Tailscale goes through: NoState  Starting  Running  Online)"
  connected=false
  consecutive_failures=0
  for i in {1..60}; do
    # Run tailscale status and capture its output so the user can see what's happening.
    # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
    ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)

    if echo "$ts_output" | grep -q "offers exit node"; then
      connected=true
      break
    fi

    # Track consecutive failures to detect a stuck state.
    # If we see 40+ consecutive 'NoState' responses, give up early
    # the daemon is alive but can't connect to the coordination server.
    if echo "$ts_output" | grep -q "NoState"; then
      consecutive_failures=$((consecutive_failures + 1))
      if [ "$consecutive_failures" -ge 40 ]; then
        echo ""
        echo "  [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
        echo "    Tailscale daemon is running but can't reach the coordination server."
        break
      fi
    else
      consecutive_failures=0
    fi

    # Print a condensed status every 10 seconds so the user can see progress
    if [ $((i % 10)) -eq 0 ]; then
      ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
      echo "  [attempt $i/60] $ts_short"
    elif [ $((i % 5)) -eq 0 ]; then
      echo -n "[$i]"
    else
      echo -n "."
    fi
    sleep 1
  done
  echo ""

  if [ "$connected" = "true" ]; then
    echo "Gateway container is online on the Tailscale mesh."
  elif [ "$BYPASS_PING" = "true" ]; then
    echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
  else
    echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
    echo "  The container was seen in the following states during the wait:" >&2
    echo "$ts_output" | sed 's/^/    /' >&2
    echo "" >&2
    echo "  This could mean:" >&2
    echo "    1. Tailscale auth key has expired  check TS_AUTHKEY in .env" >&2
    echo "    2. Network issue inside the container  check: docker compose logs tailscale" >&2
    echo "    3. gluetun VPN tunnel not connecting  check: docker compose logs warp | tail -20" >&2
    echo "" >&2
    echo "  Diagnose manually:" >&2
    echo "    docker compose exec -T tailscale tailscale status" >&2
    echo "    docker compose exec -T tailscale tailscale netcheck" >&2
    cleanup_handler ERR $LINENO
    exit 1
  fi

  # 6b. Resolve the gateway's Tailscale IP.
  # Dynamically query the container for the IP (handles IP changes after re-auth)
```

```
  # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
  #
  # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
  # returns (\r) into captured output. If $TS_IP ends with \r, then
  # `networksetup -setdnsservers` silently rejects the malformed IP and
  # DNS stays at 1.1.1.1  the primary bug this project was hitting.
  # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
  # the first field of the first line.
  TS_IP=""
  if [ "$connected" = "true" ]; then
    TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
     awk 'NR==1{print $1; exit}' || true)
    if [ -n "$TS_IP" ]; then
      echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
      echo "$TS_IP" > .gateway_ip
    else
      echo "    [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
      TS_IP=$(read_adguard_ip || true)
    fi
  else
    TS_IP=$(read_adguard_ip || true)
  fi

  # 7. Connect host Mac to Tailscale mesh.
  # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
  # system LaunchDaemon via 'brew services start tailscale'), the previous five-
  # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
  # no menu bar item to click, and no Accessibility permission required.
  #
  # CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
  # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
  # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
  # are impossible without the host on the mesh).
  #
  # We connect in two phases:
  #   Phase A  Join mesh WITHOUT exit node (so pre-flight checks can run)
  #   Phase B  Set exit node only after pre-flight checks confirm the gateway works
  HOST_ON_MESH=false
  SKIP_EXIT_NODE=false
  if [ "$HIJACK_HOST" = "false" ]; then
    echo -e "\n[Info] HIJACK_HOST is false (Headless Mode). Host networking will remain untouched."
    SKIP_EXIT_NODE=true
  elif [ -n "$TS_BIN" ]; then
    echo -n "Verifying tailscaled is reachable"
    daemon_ready=false
    status_exit=0
    for i in {1..10}; do
      # Test if the daemon responds to status check.
      # If status returns non-zero, check if it was a normal "stopped/disconnected" state
      # (exit code 1 is normal for disconnected). If the command exited quickly (not killed by timeout),
      # and the error is just "Tailscale is stopped", then the daemon is alive and ready.
      # But if it timed out (exit code 143) or is completely unresponsive, it's wedged.
      status_exit=0
      if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
        daemon_ready=true
        break
      else
        status_exit=$?
        if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
          daemon_ready=true
          break
        fi
      fi
      echo -n "."
      sleep 1
    done
    echo ""

    # If the daemon was unresponsive or socket check failed, attempt auto-restart
    if [ "$daemon_ready" != "true" ]; then
      if [ "$status_exit" -eq 143 ]; then
        warn "tailscaled daemon is wedged/unresponsive. Restarting it..."
      else
        warn "tailscaled daemon is not running. Attempting to start it..."
      fi
      restart_tailscaled_daemon

      # Re-verify
      if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
        daemon_ready=true
      fi
```

```
    fi

    if [ "$daemon_ready" = "true" ]; then
      echo "tailscaled is responsive (running as a system service)."

      #  Phase A: Join mesh without exit node
      echo "Connecting host to Tailscale mesh (no exit node yet)..."
      if $TS_BIN up; then
        $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node= || true
        HOST_ON_MESH=true
        echo "Host is on Tailscale mesh."
      else
        warn "tailscale up failed even though the daemon is running."
        echo "  This usually means the host hasn't authenticated with your tailnet yet."
        echo "  Run this command in a terminal to authenticate:"
        echo ""
        echo "      sudo tailscale up"
        echo ""
        echo "  A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
      fi
    else
      fail "tailscaled could NOT be started automatically."
      echo "  Run this in a terminal to start and authenticate Tailscale:"
      echo ""
      echo "      sudo brew services start tailscale"
      echo "      sudo tailscale up"
      echo ""
      echo "  Then re-run toggle.sh."
    fi
  else
    echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
  fi

  #  Phase B: Pre-flight checks + enable exit node only if safe
  if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
    echo -e "\n"
    echo "Pre-flight connectivity check"
    echo ""

    # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
    # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
    # ("direct connection not established") even though the pong was received.
    # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
    echo -n "  [1/3] Gateway reachable via Tailscale... "
    ping_ok=false
    # The host tailscaled needs a few seconds to propagate the new network map and
    # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
    for attempt in {1..45}; do
      if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
        ping_ok=true
        break
      fi
      sleep 1
    done
    if [ "$ping_ok" = "true" ]; then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
    # Colima's SSH tunnel only forwards TCP (not UDP), so we use +tcp.
    echo -n "  [2/3] AdGuard DNS via localhost... "
    if command -v dig &>/dev/null; then
      if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 &>/dev/null; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    elif command -v nslookup &>/dev/null; then
      if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 &>/dev/null; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    else
      echo "SKIP (dig/nslookup not available)"
    fi
```

```
2096
2097      # Check 3: Does the WARP container have external internet?
2098      echo -n "  [3/3] WARP container internet... "
2099      tmp_err=$(mktemp)
2100      if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace >/dev/null 2>
          "$tmp_err"; then
2101        echo "PASS"
2102      else
2103        echo "FAIL"
2104        if [ -s "$tmp_err" ]; then
2105          err_msg=$(cat "$tmp_err" | tr -d '\r')
2106          echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Check 3] WARP container check failed: $err_msg" >> output.log
2107        fi
2108        SKIP_EXIT_NODE=true
2109      fi
2110      rm -f "$tmp_err"
2111
2112      # Act based on check results
2113      if [ "$SKIP_EXIT_NODE" != "true" ]; then
2114        echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
2115        $TS_BIN up || true
2116        if $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node="$TS_IP" --exit-node-allow-
           lan-access=true; then
2117          echo "Exit node enabled."
2118          add_warp_bypass_routes
2119          setup_exit_node_routing
2120          enable_killswitch
2121          # Opt-in FaceTime split-route (no-op unless FACETIME_SPLIT_ROUTE=true).
2122          # After the kill-switch so the com.apple/nullexit anchor + table exist.
2123          add_apple_split_routes
2124        else
2125          echo "[Warning] Failed to set exit node."
2126          SKIP_EXIT_NODE=true
2127        fi
2128      fi
2129
2130      if [ "$SKIP_EXIT_NODE" = "true" ]; then
2131        warn "Pre-flight checks failed  EXIT NODE + DNS HIJACK SKIPPED."
2132        echo "  Host stays on Tailscale mesh but your internet routes through your normal connection."
2133        echo "  Troubleshoot with:"
2134        echo "    docker compose logs warp | tail -20"
2135        echo "    docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"
2136        echo ""
2137        echo "  Falling back to SOCKS5 proxy for traffic routing..."
2138      fi
2139    elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
2140      echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
2141      SKIP_EXIT_NODE=true
2142    fi
2143
2144    # 8. Apply DNS Hijacking.
2145    # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
2146    # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
2147    # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
2148    if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
2149      echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
2150      echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
2151      if force_dns_to_gateway "$TS_IP"; then
2152        echo "DNS hijack verified: single server = $TS_IP."
2153      else
2154        echo "[Warning] DNS hijack could not be verified."
2155        echo "  If DNS is broken, run manually:"
2156        echo "    networksetup -setdnsservers \"$ACTIVE_SERVICE\" $TS_IP"
2157        echo "    networksetup -setsearchdomains \"$ACTIVE_SERVICE\" ts.net"
2158      fi
2159    elif [ -n "$TS_IP" ] && [ "$HIJACK_HOST" != "false" ]; then
2160      echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
2161      echo "  Trying local DNS proxy for ad-blocking..."
2162      if start_local_dns_proxy; then
2163        echo "  Ad-blocking active via local DNS proxy through AdGuard."
2164      else
2165        echo "  Local DNS proxy unavailable. DNS stays at 1.1.1.1."
2166        echo "  AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
2167      fi
2168
2169      # Enable SOCKS5 proxy for traffic routing through WARP
2170      # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
2171      echo -e "\n  Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
2172      echo "  (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
2173      enable_socks_proxy "$ACTIVE_SERVICE"
2174
```

```
    # Verify the SOCKS5 proxy works through WARP
    echo -n "  Verifying SOCKS5 proxy routes through WARP... "
    if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null | grep -q "warp=on"; then
      echo "ok (warp=on)."
      echo "  All TCP traffic now encrypted through Cloudflare WARP tunnel."
    else
      echo "FAILED (proxy not routing through WARP)."
      echo "  Disabling SOCKS5 proxy  internet stays on direct connection."
      disable_socks_proxy
      echo "  SOCKS proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
    fi
  else
    echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
  fi

  # 9. Verify host connectivity
  if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
    if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
      echo "Host is online on the Tailscale mesh."
    else
      echo "[Warning] Host is not responding on Tailscale."
    fi
  fi

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."

  # Prevent system from going to sleep while gateway is running
  start_sleep_prevention

  if [[ "$HIJACK_HOST" != "false" ]]; then
    if [ -n "$TS_IP" ] && [ "$TS_IP" != "1.1.1.1" ]; then
      start_dns_watcher "$TS_IP"
    fi

    # Start WARP liveness monitor (logs to output.log on warp flip)
    start_warp_watcher
  fi

  # Reset sharing services after network configuration to prevent AirDrop freezes
  echo "Resetting macOS sharing services (AirDrop/AirPlay)..."
  reset_sharing_services

  # Tell external watchers (post-wake / network-change) the gateway is up.
  write_gateway_active_marker

  #  Local Network Surveillance Check
  echo -e "\n"
  echo "Local Network Surveillance Check"
  echo ""
  # Count populated ARP cache entries to detect network scanning or high congestion
  # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
  ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
  if [ "$ARP_COUNT" -gt 15 ]; then
    warn "High ARP activity detected ($ARP_COUNT devices in cache)."
    warn "This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)."
    echo -e "  [] Your traffic remains fully encrypted and invisible.\n"
  else
    echo -e "  [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
  fi
fi

SUCCESS_RUN=true

# Final DNS state summary
if [ -n "$ACTIVE_SERVICE" ]; then
  FINAL_DNS=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log || true)
  echo -e "\n"
  echo -e "DNS STATE: $FINAL_DNS"
  echo -e ""
fi

if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
  if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
    echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
    echo "  The gateway gracefully fell back to yesterday's cached blocklist."
    echo "  (Run 'docker logs rule-compiler' later to investigate which link died.)"
  fi
fi

```

```
rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
if [ -t 0 ]; then
  read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
  if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
    echo "Reversing gateway state..."
    exec bash "$0"
  fi
fi

echo -e "\nYou can close this terminal window now."
```

7 recover.sh

```
#!/bin/bash
# recover.sh  nullexit KILL-SWITCH recovery script (macOS + Linux)
# Fixes internet when toggle.sh leaves you stranded.
#
# This is a NUCLEAR option: it undoes EVERYTHING toggle.sh does by:
#   1. Disconnecting host Tailscale from the mesh (exit-node off)
#   2. Resetting DNS to default on ALL network services
#   3. Disabling rogue proxy settings (SOCKS, web, secure web)
#   4. Flushing the DNS cache
#   5. Flushing stale routing table entries
#   6. Stopping gateway Docker containers
#   7. Power-cycling Wi-Fi
#   8. Verifying internet is working
#
# One file, both OSes: the dispatch below runs the self-contained Linux
# recovery and exits; on macOS it falls through to the richer dual-mode
# implementation (default teardown / --post-wake refresh).
#
# Usage:  double-click Recover-Gateway.applescript
#         OR  ./recover.sh
#         OR  ./recover.sh --post-wake    (macOS only  lighter refresh, keeps gateway live)

set -e

#  Resolve script directory (used for path-relative refs)
# recover.sh lives at the repo root; SCRIPT_DIR lets us launch toggle.sh
# (and reference any other repo-root file) from the same directory no
# matter where the user invoked recover.sh from.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
cd "$SCRIPT_DIR" || exit 1

#
# LINUX RECOVERY  self-contained; runs the full teardown then exits before the
# macOS body below. Native Linux tooling (resolvectl / nmcli / ip) throughout.
#
if [[ "$OSTYPE" != "darwin"* ]]; then
  # Source common formatting and helper functions
  source "$SCRIPT_DIR/scripts/common.sh"

  # Always add Homebrew path
  export PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:$PATH"

  echo ""
  echo -e "${RED}${BOLD}${NC}"
  echo -e "${RED}${BOLD}     Nullexit Gateway  RECOVERY MODE     ${NC}"
  echo -e "${RED}${BOLD}${NC}"
  echo ""
  echo "This script will tear down the gateway and restore normal internet."
  echo ""

  #  Cache sudo credentials upfront
  # Prevents the script from hanging halfway through when it needs to run
  # privileged commands (route flush, Wi-Fi power cycle, etc.).
  echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
  sudo -v
  (
    while true; do
      sudo -n true 2>/dev/null
      sleep 60
    done
  ) &
  SUDO_KEEPER_PID=$!
  trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
```

```
64
65  #  1. Disconnect host Tailscale
66  step "Disconnecting host Tailscale from mesh"
67  if command -v tailscale >> output.log 2>&1; then
68    # Check if actually connected first (with timeout  tailscaled could be wedged)
69    if run_with_timeout 5 tailscale status >> output.log 2>&1; then
70      # Also explicitly reset any exit-node so it doesn't linger.
71      if run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= >> output.log
      2>&1; then
72        ok "Exit-node preference cleared"
73      else
74        warn "tailscale up --reset didn't respond (tailscaled may be wedged)"
75      fi
76      ts_args=""
77      if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
78      if run_with_timeout 10 tailscale down $ts_args >> output.log 2>&1; then
79        ok "Tailscale disconnected"
80      else
81        warn "tailscale down didn't respond"
82      fi
83    else
84      warn "tailscaled not reachable  skipping (timeout after 5s)"
85    fi
86  else
87    warn "tailscale CLI not found  skipping"
88  fi
89
90  #  2. Reset DNS to default on ALL network services
91  step "Resetting DNS to default on active interface"
92
93  ACTIVE_SERVICE=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
94  if [ -n "$ACTIVE_SERVICE" ]; then
95    sudo resolvectl revert "$ACTIVE_SERVICE" 2>> output.log || true
96    ok "DNS reset on $ACTIVE_SERVICE"
97  else
98    warn "Could not determine active network interface"
99  fi
100
101 #  3. Disable rogue proxy settings
102 step "Disabling proxy settings (SOCKS, web, secure web)"
103 echo "  (Linux global SOCKS proxy configuration skipped.)"
104 ok "Proxy settings cleared"
105
106
107 #  4. Flush DNS cache
108 step "Flushing DNS cache"
109 if command -v resolvectl >> output.log 2>&1; then
110   sudo resolvectl flush-caches 2>> output.log || true
111   ok "DNS cache flushed"
112 else
113   warn "resolvectl not found"
114 fi
115
116 #  5. Clear stale routing table
117 step "Flushing stale routing table entries"
118 sudo ip route flush cache >> output.log 2>&1 || true
119 ok "Routing table flushed"
120
121 #  6. Stop gateway Docker containers
122 step "Stopping gateway Docker containers"
123 if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
124   if [ -f "docker-compose.yml" ]; then
125     docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
126   else
127     warn "docker-compose.yml not found in current directory"
128   fi
129 else
130   warn "Docker not available  skipping"
131 fi
132
133 #  6b. Stopping sleep prevention (systemd-inhibit)
134 step "Stopping sleep prevention"
135 PID_FILE="/tmp/nullexit-caffeinate.pid"
136 if [ -f "$PID_FILE" ]; then
137   INHIBIT_PID=$(cat "$PID_FILE")
138   if [ -n "$INHIBIT_PID" ] && kill -0 "$INHIBIT_PID" 2>/dev/null; then
139     kill "$INHIBIT_PID" 2>/dev/null || true
140     ok "Sleep prevention stopped (PID $INHIBIT_PID)"
141   else
142     warn "Sleep prevention process not running, cleaning up stale PID file"
143   fi
```

```
    rm -f "$PID_FILE"
  else
    ok "Sleep prevention was not active"
  fi

  #  7. Power-cycle Wi-Fi
  step "Power-cycling Wi-Fi"
  if command -v nmcli &>/dev/null; then
    sudo nmcli radio wifi off 2>> output.log || true
    sleep 2
    sudo nmcli radio wifi on 2>> output.log || true
    ok "Wi-Fi bounced via nmcli"
    sleep 3
  else
    warn "nmcli not found. Falling back to ip link..."
    WIFI_IFACE=$(ip link | grep -E '^[0-9]+: (wl[a-zA-Z0-9]+|wlan[0-9]+):' | awk -F': ' '{print $2}')
    if [ -n "$WIFI_IFACE" ]; then
      sudo ip link set "$WIFI_IFACE" down 2>> output.log || true
      sleep 2
      sudo ip link set "$WIFI_IFACE" up 2>> output.log || true
      ok "Wi-Fi bounced via ip link (interface $WIFI_IFACE)"
      sleep 3
    else
      warn "No Wi-Fi interface detected."
    fi
  fi

  #  8. Verify internet connectivity
  verify_internet_connectivity

  #  Summary
  echo ""
  echo -e "${BOLD}${NC}"
  echo -e "${BOLD}RECOVERY SUMMARY${NC}"
  echo -e "${BOLD}${NC}"

  # Show final DNS state
  FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr
     '\n' ' ' || echo "N/A")
  echo "  DNS ($ACTIVE_SERVICE):  $FINAL_DNS"

  echo ""

  if [ "$INTERNET_OK" = "true" ]; then
    echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
  else
    echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
    echo "    This might mean:"
    echo "    - Your Wi-Fi is disconnected from the router"
    echo "    - Tailscale is still interfering (try restarting tailscaled:"
    echo "        systemctl restart tailscaled)"
    echo "    - You need to re-connect to Wi-Fi in the menu bar"
    echo "    - Try toggling Wi-Fi off/on in the menu bar"
    echo ""
    echo "    Or try manually:"
    echo "      resolvectl revert $ACTIVE_SERVICE"
    extra_args=""
    if is_kill_switch_enabled; then extra_args="--accept-risk=lose-ssh"; fi
    echo "      tailscale down $extra_args"
    echo "      sudo ip route flush cache"
    echo "      systemctl restart tailscaled"
  fi

  echo ""
  echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
  echo ""

  exit 0
fi

#
# macOS RECOVERY (dual-mode: default teardown / --post-wake refresh)
#

# Enforce Cryptographic Script Integrity
if [ -f "scripts/crypto.sh" ]; then
  if ! bash scripts/crypto.sh --verify; then
    exit 1
  fi
fi

```

```
#  Argument parsing
# --post-wake: lightweight refresh after sleep/wake or Wi-Fi roam.
#              Keeps the gateway live while HUPing DNS caches, refreshing the
#              Tailscale exit-node leak, re-hijacking host DNS, and force-
#              recreating the warp container if its gluetun healthcheck failed.
#              Does NOT tear down Docker, Tailscale, sleep prevention, or Wi-Fi.
#              Called from scripts/watcher.sh on every wake / network-change event
#              while /tmp/nullexit-gateway-active.marker is present (i.e. the
#              gateway is currently up). See devref.md 10.29 for the why.
POST_WAKE=false
AUTO_MODE=false
for arg in "$@"; do
  if [ "$arg" = "--post-wake" ]; then
    POST_WAKE=true
  elif [ "$arg" = "--auto" ] || [ "$arg" = "--non-interactive" ]; then
    AUTO_MODE=true
  fi
done

rm -f "$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker"

# Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
LOCK_FILE="/tmp/nullexit-toggle.lock"
if [ -f "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
      if [ "$POST_WAKE" = "true" ]; then
        echo "[$(date -u +%FT%TZ)] POST-WAKE skipped: another lifecycle script (PID $LOCK_PID) is running." >>
    output.log
        exit 0
      else
        echo "Error: Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running." >&2
        exit 1
      fi
    fi
  fi
fi
echo "$$" > "$LOCK_FILE"

# Clear the active-state marker in nuclear recovery mode since the gateway is being
# completely torn down. This prevents scripts/watcher.sh from triggering post-wake
# recoveries in response to network changes caused by the teardown itself.
if [ "$POST_WAKE" = "false" ]; then
  rm -f "/tmp/nullexit-gateway-active.marker"
fi

# WARP Watcher inhibit marker: written immediately in --post-wake mode so the
# background WARP Watcher (in toggle.sh) skips failure counting while we are
# intentionally tearing down / recreating the warp container. Without this, the
# watcher accumulates 6 consecutive warp=off readings during force-recreate (~35s)
# and fires nuclear recover.sh, killing a gateway that would have been healthy.
WARP_INHIBIT_FILE="/tmp/nullexit-warp-inhibit.marker"
if [ "$POST_WAKE" = "true" ]; then
  touch "$WARP_INHIBIT_FILE"
  echo "[$(date -u +%FT%TZ)] POST-WAKE started  WARP Watcher inhibited to prevent spurious shutdown during warp
      recreate." >> output.log
fi

# Ensure the inhibit marker and lock file are always removed when the script exits (success or error)
cleanup_recover() {
  rm -f "$WARP_INHIBIT_FILE"
  if [ -f "$LOCK_FILE" ] && [ "$(cat "$LOCK_FILE" 2>/dev/null)" = "$$" ]; then
    rm -f "$LOCK_FILE"
  fi
}
trap cleanup_recover EXIT

# Source common formatting and helper functions
source "$SCRIPT_DIR/scripts/common.sh"

# Always add Homebrew path
setup_standard_path

echo ""
if [ "$POST_WAKE" = "true" ]; then
  echo -e "${YELLOW}${BOLD}${NC}"
  echo -e "${YELLOW}${BOLD}  Nullexit  POST-WAKE / POST-ROAM LIGHT ${NC}"
  echo -e "${YELLOW}${BOLD}  (keeps the gateway live, refreshes)    ${NC}"
  echo -e "${YELLOW}${BOLD}${NC}"
  echo ""
```

```
  echo "Re-hijacking DNS, refreshing Tailscale exit-node, force-recreating"
  echo "warp if unhealthy, and resetting sharing services. The gateway stays"
  echo "up. docker compose / Wi-Fi / caffeinate are NOT touched."
  echo ""
else
  echo -e "${RED}${BOLD}${NC}"
  echo -e "${RED}${BOLD}     Nullexit Gateway  RECOVERY MODE     ${NC}"
  echo -e "${RED}${BOLD}${NC}"
  echo ""
  echo "This script will tear down the gateway and restore normal internet."
  echo ""
fi

#  Cache sudo credentials upfront (default mode ONLY)
# Skip in --post-wake: the LaunchAgent-launched watcher has no TTY, so `sudo -v`
# would block forever waiting for a password. All post-wake commands use `sudo -n`
# which is non-blocking and relies on /etc/sudoers.d/nullexit NOPASSWD entries.
SUDO_KEEPER_PID=""
if [ "$POST_WAKE" = "false" ] && [ "$AUTO_MODE" = "false" ] && [ -t 0 ]; then
  echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
  sudo -v
  (
    while true; do
      sudo -n true 2>/dev/null
      sleep 60
    done
  ) &
  SUDO_KEEPER_PID=$!
  trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
fi

#  Resolve physical default interface and gateway upfront
PHYSICAL_IFACE=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
    getline; print $2; exit}')
[ -z "$PHYSICAL_IFACE" ] && PHYSICAL_IFACE="en0"

if [ "$POST_WAKE" = "true" ]; then
  wait_for_dhcp_settle
else
  # Quick resolution in non-post-wake mode (non-blocking)
  PHYSICAL_GW=$(ipconfig getpacket "$PHYSICAL_IFACE" 2>> output.log | awk -F'[{}]' '/router /{print $2}')
fi

#  1. Disconnect host Tailscale (default) / Refresh exit-node (post-wake)

if [ "$POST_WAKE" = "true" ]; then
  step "Refreshing host Tailscale exit-node routing (post-wake)"
  if command -v tailscale >> output.log 2>&1; then
    # Re-issue the SAME exit-node up command toggle.sh START uses. This refreshes
    # the DERP relay mapping and re-asserts the exit-node preference without
    # dropping the host's mesh connection (which `tailscale down` would).
    TS_IP=""
    for src in .gateway_ip; do
      [ -f "$src" ] || continue
      TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
      [ -n "$TS_IP" ] && break
    done

    if [ -n "$TS_IP" ]; then
      step "Re-establishing exit node $TS_IP via 'tailscale set --exit-node'"
      # `tailscale set` only mutates the named preference (exit-node here).
      # Crucially it does NOT call --reset, which would re-apply ALL flags and
      # trigger a sub-second route-table re-evaluation cycle each post-wake.
      # On already-connected nodes, the lighter `set` keeps the existing tunnel
      # alive while only changing the exit-node preference.
      if run_with_timeout 15 tailscale set --exit-node="$TS_IP" \
                            --exit-node-allow-lan-access=true \
                            >> output.log 2>&1; then
        ok "Exit-node $TS_IP re-asserted via 'tailscale set'"
      else
        warn "tailscale set --exit-node=$TS_IP didn't respond; falling back to mesh-only"
        run_with_timeout 10 tailscale set --exit-node= \
                          >> output.log 2>&1 || true
      fi
    else
      warn ".gateway_ip missing  cannot re-assert exit node"
      run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= \
                          >> output.log 2>&1 || true
    fi
  else
    warn "tailscale CLI not found"
```

```
  fi
else
  step "Disconnecting host Tailscale from mesh"
  disconnect_tailscale_host "tailscale"
fi

#  2. Reset DNS to default on ALL network services (default) / Re-hijack gateway DNS (post-wake)
if [ "$POST_WAKE" = "true" ]; then
  step "Re-hijacking host DNS to gateway IP (post-wake / post-roam)"
  TS_IP=""
  for src in .gateway_ip; do
    [ -f "$src" ] || continue
    TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
    [ -n "$TS_IP" ] && break
  done

  if [ -n "$TS_IP" ]; then
    # Detect the active network service (using helper in common.sh)
    ACTIVE_SVC=$(get_active_service)
    EN0_SVC=$(get_en0_service)

    networksetup -setsearchdomains "$ACTIVE_SVC" "ts.net" 2>> output.log || true
    networksetup -setdnsservers "$ACTIVE_SVC" "$TS_IP" 2>> output.log || true
    networksetup -setsearchdomains "$EN0_SVC" "ts.net" 2>> output.log || true
    networksetup -setdnsservers "$EN0_SVC" "$TS_IP" 2>> output.log || true

    HITS=$(networksetup -getdnsservers "$ACTIVE_SVC" 2>> output.log | grep -Fx "$TS_IP" || true)
    if [ -n "$HITS" ]; then
      ok "DNS re-hijacked to $TS_IP on services: $ACTIVE_SVC, $EN0_SVC"
    else
      warn "networksetup accepted but DNS read-back didn't include $TS_IP"
    fi
  else
    warn ".gateway_ip missing  cannot re-hijack DNS (user must run toggle.sh START)"
  fi
else
  step "Resetting DNS to default on all network services (empty = DHCP)"

  # Get ALL network services (handles spaces in names, skips disabled ones)
  SERVICES=$(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
    true)

  if [ -z "$SERVICES" ]; then
    # Fallback to common names
    SERVICES="Wi-Fi Ethernet Thunderbolt Ethernet USB 10/100 LAN"
  fi

  RESET_COUNT=0
  while IFS= read -r service; do
    # Strip leading asterisk (disabled services)
    service_clean=$(echo "$service" | sed 's/^\*//')
    [ -z "$service_clean" ] && continue

    # Check if this service has a hardware port (skip VPN/service-only entries)
    if networksetup -listallhardwareports 2>> output.log | grep -A1 "Port: $service_clean$" | grep -q "Device:"
      || [ "$service_clean" = "Wi-Fi" ]; then
      (
        networksetup -setsearchdomains "$service_clean" "Empty" 2>> output.log
        # Set to "empty" so macOS uses DHCP-assigned DNS (not hardcoded 1.1.1.1)
        sudo networksetup -setdnsservers "$service_clean" "empty" 2>> output.log
      ) && RESET_COUNT=$((RESET_COUNT + 1)) || true
    fi
  done <<< "$SERVICES"

  ok "DNS reset on $RESET_COUNT service(s)"
fi

#  3. Disable rogue proxy settings (default only  gateway uses proxy in non-exit-node path)
if [ "$POST_WAKE" = "false" ]; then
  step "Disabling proxy settings (SOCKS, web, secure web)"
  for svc in $(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
    echo "Wi-Fi"); do
    svc_clean=$(echo "$svc" | sed 's/^\*//')
    [ -z "$svc_clean" ] && continue
    disable_all_proxies "$svc_clean"
  done
  ok "Proxy settings cleared"
else
  step "Proxy settings: leaving untouched (post-wake keeps gateway SOCKS wiring alive)"
fi

```

```
#  4. Flush macOS DNS cache (always  safe and useful in both modes)
step "Flushing DNS cache"
if command -v dscacheutil >> output.log 2>&1; then
  flush_dns_cache
  ok "DNS cache flushed (mDNSResponder HUP)"
else
  warn "dscacheutil not found"
fi

#  5. Clear stale routing table (always  safe in both modes)
step "Flushing stale routing table entries"
# Resolve physical default gateway IP before flushing
# We cannot trust `route get default` here because Tailscale may have already hijacked
# the default route to utunX. We must read the actual DHCP router assignment from the hardware.
# PHYSICAL_IFACE and PHYSICAL_GW resolved upfront

# Instead of full route -n flush (which destroys macOS loopback/multicast and wedges the OS until reboot),
# we cleanly delete the Tailscale overrides and Cloudflare WARP bypass routes in ALL modes.
# This ensures tailscaled isn't blocked from reaching its control plane when waking from sleep.
sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
remove_warp_bypass_routes
ok "Routing overrides cleared"

if [ "$POST_WAKE" = "false" ]; then
  disable_killswitch
  ok "Routing table flush completed via targeted deletes and firewall disabled"
else
  ok "Routing overrides cleared (post-wake mode to preserve Colima bridge100)"
fi

# In post-wake mode we don't bounce Wi-Fi, so we MUST manually restore the physical default route
if [ "$POST_WAKE" = "true" ] && [ -n "$PHYSICAL_GW" ]; then
  # Ensure physical gateway is set just in case it was dropped
  sudo -n route add default "$PHYSICAL_GW" >> output.log 2>&1 || true
fi

# CONDITIONAL BYPASS ROUTE RESTORATION:
# If this was a post-wake recovery and the exit node is active, the route flush
# just wiped the static bypass routes for the WARP endpoints. We must re-add
# them immediately to prevent a routing loop when Tailscale starts sending traffic.
if [ "$POST_WAKE" = "true" ]; then
  if [ -z "${ACTIVE_IF:-}" ]; then
    ACTIVE_IF=$(route get default 2>> output.log | awk '/interface:/{print $2; exit}')
    if [ -z "$ACTIVE_IF" ] || [[ "$ACTIVE_IF" =~ ^utun ]] || [[ "$ACTIVE_IF" =~ ^tun ]]; then
      for i in en0 en1 en2 en3; do
        if ifconfig "$i" 2>> output.log | grep -q 'status: active'; then
          ACTIVE_IF="$i"; break
        fi
      done
    fi
    [ -z "$ACTIVE_IF" ] && ACTIVE_IF="en0"
  fi

  echo "Re-adding host bypass routes for Cloudflare WARP endpoints..."
  remove_warp_bypass_routes
  if [ -n "${PHYSICAL_GW:-}" ]; then
    add_warp_bypass_routes "$PHYSICAL_GW"
  else
    add_warp_bypass_routes "$ACTIVE_IF"
  fi
  # Re-assert the opt-in FaceTime split-route so it survives roam/wake
  # (no-op unless FACETIME_SPLIT_ROUTE=true). Physical gateway is known here.
  add_apple_split_routes "${PHYSICAL_GW:-}"

  # Also re-apply the default route pointing to the Tailscale interface
  ts_iface=$(ifconfig 2>> output.log | grep -B4 "inet 100." | grep -E '^[a-z0-9]+' | cut -d: -f1 | head -n 1)
  if [ -n "$ts_iface" ]; then
    echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
    # DO NOT delete the physical default route! It crashes Colima.
    # Use the /1 trick to mathematically override the default route.
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    enable_killswitch
  fi
fi

#  6. Stop gateway Docker containers (default) / Force-recreate warp if unhealthy (post-wake)
if [ "$POST_WAKE" = "true" ]; then
```

```
  #  6a. Check for Roaming Network Mismatch (Bridged vs Enterprise Wi-Fi)
  if [ -f "/tmp/nullexit_colima_mode.txt" ] && [ -f "$SCRIPT_DIR/../.lan_p2p_detected" ]; then
    current_colima_mode=$(cat /tmp/nullexit_colima_mode.txt)
    p2p_allowed=$(cat "$SCRIPT_DIR/../.lan_p2p_detected")
    required_colima_mode="bridged"
    [ "$p2p_allowed" == "false" ] && required_colima_mode="shared"

    if [ "$current_colima_mode" != "$required_colima_mode" ]; then
      warn "Network roaming mismatch detected!"
      warn "Colima is running in '$current_colima_mode' mode but new network requires '$required_colima_mode'."
      warn "Executing full gateway restart to protect against MAC-spoofing deauthentication..."
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Roam] Mismatch ($current_colima_mode vs $required_colima_mode). Triggering toggle.sh --restart." >> output.log
      nohup bash "$SCRIPT_DIR/../toggle.sh" --restart >/dev/null 2>&1 &
      exit 0
    fi
  fi

  step "Checking Colima VM state"
  colima_status=$(colima status 2>&1)
  if echo "$colima_status" | grep -qi "running"; then
    ok "Colima VM is running"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima status check passed." >> output.log
  else
    warn "Colima VM is NOT running or unresponsive"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima is unhealthy or dead! Output: $colima_status" >> output.log
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Attempting to restart Colima..." >> output.log
    colima restart >> output.log 2>&1 || warn "Failed to restart Colima"
  fi

  step "Checking warp container health (gluetun UDP tunnel)"
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Inspecting warp container state..." >> output.log
  # The warp container is the most failure-prone link in the chain at wake/roam:
  # its UDP NAT binding to Cloudflare's edge (162.159.192.1:2408) silently dies
  # during a Wi-Fi roam. We verify the tunnel is actually passing traffic via curl.
  WARP_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container status: $WARP_STATE" >> output.log

  if [ "$WARP_STATE" = "missing" ]; then
    warn "warp container is missing  leaving gateway containers alone (full relaunch requires \"$SCRIPT_DIR/toggle.sh\")"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container missing from Docker engine. Requires full toggle.sh launch." >> output.log
  else
    tmp_err=$(mktemp)
    if docker compose exec -T warp wget -qO- --timeout=3 https://www.cloudflare.com/cdn-cgi/trace 2>"$tmp_err" | grep -q "warp=on"; then
      ok "warp container is healthy and actively tunneling traffic (no recreate needed)"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container passed live traffic healthcheck (Cloudflare trace successful)." >> output.log
    else
      warn "warp container tunnel is dead (Docker status: $WARP_STATE)  force-recreating all gateway containers to restore network namespace"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container is DEAD or STUCK. Forcing recreation of gateway stack..." >> output.log
      if [ -s "$tmp_err" ]; then
        err_msg=$(cat "$tmp_err" | tr -d '\r')
        echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Healthcheck failure details: $err_msg" >> output.log
      fi
      docker compose up -d --force-recreate >> output.log 2>&1 || warn "force-recreate failed"

      NEW_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
      ok "warp container health after recreate: $NEW_STATE"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container health after recreation: $NEW_STATE" >> output.log
    fi
    rm -f "$tmp_err"
  fi
else
  step "Stopping gateway Docker containers"
  if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
    if [ -f "docker-compose.yml" ]; then
      docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
    else
      warn "docker-compose.yml not found in current directory"
    fi
  else
    warn "Docker not available  skipping"
  fi
```

```
fi

#  6b. Stopping sleep prevention (caffeinate)  default only
if [ "$POST_WAKE" = "false" ]; then
  step "Stopping sleep prevention"
  stop_pidfile_daemon "$PID_CAFFEINATE" "system sleep prevention"
else
  step "Sleep prevention: leaving untouched (caffeinate already re-acquired by parent shell)"
fi

#  6c. Stopping DNS Watcher  default only
if [ "$POST_WAKE" = "false" ]; then
  step "Stopping DNS Watcher"
  stop_pidfile_daemon "$PID_DNS_WATCHER" "background DNS Watcher"
else
  step "DNS Watcher: leaving untouched (it keeps re-hijacking DNS on every roam)"
fi

#  7. Power-cycle Wi-Fi  default only
if [ "$POST_WAKE" = "false" ]; then
  step "Power-cycling Wi-Fi"
  bounce_wifi_interfaces
else
  step "Wi-Fi: leaving untouched (post-wake keeps the current Wi-Fi association)"
fi

#  8. Verify (default checks direct internet; post-wake verifies the gateway)
if [ "$POST_WAKE" = "true" ]; then
  step "Verifying gateway is healthy after post-wake refresh"

  # Wait a moment for tailscale + warp healthcheck to settle
  sleep 3

  GATEWAY_OK=false
  if command -v dig >> output.log 2>&1; then
    # Source procedural ports if available
    if [ -f "/tmp/nullexit-ports.env" ]; then
      source "/tmp/nullexit-ports.env"
    fi
    DNS_PORT=${DNS_PROXY_PORT:-5354}

    # AdGuard is exposed at localhost:${DNS_PORT}. If this resolves, the warp tunnel
    # is functioning properly because AdGuard routes upstream to warp.
    if dig +tcp @127.0.0.1 -p "$DNS_PORT" google.com +short +timeout=5 >> output.log 2>&1; then
      ok "Gateway DNS (AdGuard via localhost:${DNS_PORT}) works"
      GATEWAY_OK=true
    else
      warn "Gateway DNS not responding yet"
    fi
  fi

  # Direct host  gateway check via Tailscale ping (relayed pong counts as success)
  if command -v tailscale >> output.log 2>&1; then
    TS_IP=""
    [ -f .gateway_ip ] && TS_IP=$(read_adguard_ip || true)
    if [ -n "$TS_IP" ] && tailscale ping --until-direct=false -c 1 --timeout 4s "$TS_IP" 2>> output.log | grep
     -q pong; then
      ok "Gateway $TS_IP reachable via Tailscale"
      GATEWAY_OK=true
    else
      warn "Tailscale ping to gateway did not pong  exit-node may still be re-establishing"
    fi
  fi

  if [ "$GATEWAY_OK" = "true" ]; then
    echo -e "  ${GREEN}${BOLD} Gateway is healthy after post-wake refresh.${NC}"
  else
    echo -e "  ${YELLOW}${BOLD} Gateway recovery is in-progress.${NC}"
    echo "    The route may need another 30s before queries succeed."
    echo "    If it stays broken for >60s, run: bash \"$SCRIPT_DIR/toggle.sh\""
  fi
else
  verify_internet_connectivity
fi

#  Summary
echo ""
echo -e "${BOLD}${NC}"
if [ "$POST_WAKE" = "true" ]; then
  echo -e "${BOLD}POST-WAKE SUMMARY${NC}"
else
```

```
  echo -e "${BOLD}RECOVERY SUMMARY${NC}"
fi
echo -e "${BOLD}${NC}"

# Show final DNS state
FINAL_DNS=$(networksetup -getdnsservers "Wi-Fi" 2>/dev/null || echo "N/A")
echo "  DNS (Wi-Fi):  $FINAL_DNS"

FINAL_DNS2=$(networksetup -getdnsservers "Ethernet" 2>/dev/null || true)
if [[ -z "$FINAL_DNS2" || "$FINAL_DNS2" == *"not a recognized"* || "$FINAL_DNS2" == *"Error"* ]]; then
  FINAL_DNS2="N/A (Not configured)"
fi
echo "  DNS (Ethernet): $FINAL_DNS2"

echo ""

if [ "$POST_WAKE" = "false" ]; then
  if [ "$INTERNET_OK" = "true" ]; then
    echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
  else
    echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
    echo "    This might mean:"
    echo "    - Your Wi-Fi is disconnected from the router"
    echo "    - Tailscale is still interfering (try restarting tailscaled:"
    echo "        brew services restart tailscale)"
    echo "    - You need to re-connect to Wi-Fi in the menu bar"
    echo "    - Try toggling Wi-Fi off/on in the menu bar"
    echo ""
    echo "    Or try manually:"
    echo "      networksetup -setdnsservers Wi-Fi empty"
    echo "      tailscale down"
    echo "      sudo route delete -net 0.0.0.0/1"
    echo "      sudo route delete -net 128.0.0.0/1"
    echo "      brew services restart tailscale"
  fi
fi

# Reset sharing services to prevent AirDrop freezes after interface changes
# This is the SAME step in BOTH modes (post-wake inherits the previous fix from 10.27).
echo -e "  Resetting macOS sharing services (AirDrop/AirPlay)..."
reset_sharing_services

echo ""
if [ "$POST_WAKE" = "true" ]; then
  echo -e "${YELLOW}${BOLD}Post-wake refresh complete.${NC}"
  # Remove the WARP Watcher inhibit marker now that the gateway is fully healthy.
  # The watcher will resume normal liveness monitoring on its next 5s poll.
  rm -f "$WARP_INHIBIT_FILE"
  echo "[$(date -u +%FT%TZ)] POST-WAKE complete  WARP Watcher inhibit released." >> output.log
else
  echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
fi
echo ""
```

8 setup.sh

```
#!/bin/bash
# setup.sh  nullexit one-shot setup script (macOS)
# Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
set -euo pipefail

# Source common formatting and helper functions
source "$(dirname "$0")/scripts/common.sh"

# Run common installation logic
source "$(dirname "$0")/scripts/setup-common.sh"

#  Compile macOS Shortcuts
echo -e "\n${BOLD}Compiling macOS Desktop Shortcuts...${NC}"
if command -v osacompile &> /dev/null; then
    osacompile -o "Toggle Gateway.app" "Toggle-Gateway.applescript" >/dev/null 2>&1
    osacompile -o "Recover Gateway.app" "Recover-Gateway.applescript" >/dev/null 2>&1
    echo "   Created 'Toggle Gateway.app' and 'Recover Gateway.app'"
fi

#  Done
```

```
echo ""
echo -e "${GREEN}${BOLD}${NC}"
echo -e "${GREEN}${BOLD}            Setup complete!                  ${NC}"
echo -e "${GREEN}${BOLD}${NC}"
echo ""
echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
echo "   https://login.tailscale.com/admin/machines"
echo "   Find your gateway  '...'  'Edit route settings'  enable Exit Node"
echo ""

echo -e "${YELLOW}${BOLD}Sudoers / Background Execution Requirements:${NC}"
echo "  To allow silent, non-interactive execution of the firewall and network routes (especially during sleep/wake recovery), you must configure passwordless sudo."
echo "  Run this exact command in your terminal (grants are scoped to least privilege):"
echo "    echo \"\$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /sbin/route, /sbin/ifconfig, /usr/bin/dscacheutil -flushcache, /usr/bin/killall -HUP mDNSResponder, /usr/bin/killall sharingd rapportd, /usr/bin/pkill -9 tailscaled, /usr/bin/pkill -f dns-proxy.py, /bin/kill, /usr/bin/true, /opt/homebrew/bin/brew services * tailscale, /opt/homebrew/bin/brew uninstall tailscale, /usr/local/bin/brew services * tailscale, /usr/local/bin/brew uninstall tailscale, /usr/bin/python3 -I \$PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I \$PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I \$PWD/scripts/dns-proxy.py\" | sudo tee /etc/sudoers.d/nullexit >/dev/null && sudo visudo -cf /etc/sudoers.d/nullexit"
echo ""

if [[ -n "${TS_IP:-}" ]]; then
    echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
    echo ""
fi

echo -e "${BOLD}To toggle the gateway on/off:${NC}"
echo "  macOS:   double-click the 'Toggle Gateway' app icon!"
echo "  Windows: double-click Toggle-Gateway.bat"
echo ""

if [[ "$OSTYPE" == "darwin"* ]]; then
    #  Install Taildrop auto-receive watcher
    echo -e "\n${BOLD}Installing Taildrop auto-receive watcher...${NC}"
    mkdir -p "$HOME/Library/LaunchAgents"
    NULLEXIT_HOME="$(cd "$(dirname "$0")" && pwd)"
    TAILDROP_PLIST="$HOME/Library/LaunchAgents/com.nullexit.taildrop-watch.plist"
    launchctl unload "$TAILDROP_PLIST" >/dev/null 2>&1 || true
    cp "$NULLEXIT_HOME/launchd/com.nullexit.taildrop-watch.plist" "$TAILDROP_PLIST"
    sed -i '' "s|__NULLEXIT_HOME__|$NULLEXIT_HOME|" "$TAILDROP_PLIST"
    launchctl load -w "$TAILDROP_PLIST"
    echo "   Incoming Taildrop files now land in ~/Downloads automatically."

    # LuLu localhost filtering check
    if [ -d "/Applications/LuLu.app" ] || systemextensionsctl list 2>/dev/null | grep -q "com.objective-see.lulu"; then
        echo -e "${YELLOW}${BOLD}LuLu Firewall Detected:${NC}"
        echo "  LuLu's localhost (loopback) filtering is OFF by default. To properly protect the"
        echo "  local nullexit proxy ports from malicious local processes, you must manually enable"
        echo "  'Allow Loopback' (or block it selectively) in LuLu's Preferences -> Rules."
        echo ""
    fi

    echo -e "${YELLOW}${BOLD}Note on macOS Tailscale Sandboxing:${NC}"
    echo "  The Tailscale GUI app (tailscale-app) installed on macOS is sandboxed."
    echo "  While you cannot run a Tailscale SSH server on this Mac host, you can"
    echo "  still SSH outbound to other devices on the mesh using their MagicDNS"
    echo "  addresses directly (e.g. ssh user@device.tailnet-name.ts.net)."
    echo ""
fi
```

9 scripts/setup-linux.sh

```
#!/bin/bash
# scripts/setup-linux.sh  nullexit one-shot setup script (Linux)
# Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
set -euo pipefail

# Source common formatting and helper functions
source "$(dirname "$0")/common.sh"

# Run common installation logic
source "$(dirname "$0")/setup-common.sh"

```

```
#  Compile Linux Desktop Shortcuts
echo -e "\n${BOLD}Creating Linux Desktop Shortcuts...${NC}"
DIR="$(pwd)"
cat <<EOF > "Toggle Gateway.desktop"
[Desktop Entry]
Version=1.0
Name=Toggle Gateway
Comment=Start or Stop Nullexit Gateway
Exec=bash -c "cd '$DIR' && ./toggle.sh; read -p 'Press Enter to close...' dummy"
Icon=utilities-terminal
Terminal=true
Type=Application
Categories=Network;Security;
EOF

cat <<EOF > "Recover Gateway.desktop"
[Desktop Entry]
Version=1.0
Name=Recover Gateway
Comment=Emergency DNS and Network Recovery
Exec=bash -c "cd '$DIR' && ./recover.sh; read -p 'Press Enter to close...' dummy"
Icon=utilities-terminal
Terminal=true
Type=Application
Categories=Network;Security;
EOF
echo "   Created 'Toggle Gateway.desktop' and 'Recover Gateway.desktop'"

#  Done
echo ""
echo -e "${GREEN}${BOLD}${NC}"
echo -e "${GREEN}${BOLD}            Setup complete!                  ${NC}"
echo -e "${GREEN}${BOLD}${NC}"
echo ""
echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
echo "   https://login.tailscale.com/admin/machines"
echo "   Find your gateway  '...'  'Edit route settings'  enable Exit Node"
echo ""

if [[ -n "${TS_IP:-}" ]]; then
    echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
    echo ""
fi

echo -e "${BOLD}To toggle the gateway on/off:${NC}"
echo "  Linux: double-click the 'Toggle Gateway' desktop icon!"
echo "  (or run ./toggle.sh in terminal)"
echo ""
```

10 docker-compose.yml

```
services:
  warp:
    image: qmcgaw/gluetun:v3.41.1
    container_name: warp
    hostname: nullexit-gateway
    cap_add:
      - NET_ADMIN
    devices:
      - /dev/net/tun:/dev/net/tun
    ports:
      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/udp"    # AdGuard DNS: host binds 5354 (the mapped DNS-proxy
     port); 5335 is AdGuard's in-container port
      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/tcp"    # AdGuard DNS (TCP)
      - "41642:41642/udp"  # Tailscale WireGuard (direct connection to host)
      - "127.0.0.1:${SOCKS_PROXY_PORT:-1080}:1080/tcp"    # SOCKS5 proxy (routed through WARP tunnel)
      - "127.0.0.1:${TOR_SOCKS_PORT:-9050}:9050/tcp"  # Tor SOCKS5 proxy (procedurally generated)
      - "127.0.0.1:${TOR_CONTROL_PORT:-9051}:9051/tcp"  # Tor Control port (procedurally generated)
    environment:
      - TZ=UTC
      - VPN_SERVICE_PROVIDER=custom
      - VPN_TYPE=wireguard
      - WIREGUARD_PRIVATE_KEY=${WIREGUARD_PRIVATE_KEY}
      - WIREGUARD_PUBLIC_KEY=${WIREGUARD_PUBLIC_KEY}
      - WIREGUARD_ADDRESSES=${WIREGUARD_ADDRESSES}
      - WIREGUARD_ENDPOINT_IP=${WARP_ENDPOINT_1:-162.159.192.1}
```

```
      - WIREGUARD_ENDPOINT_PORT=2408
      # Set to 25s (WireGuard default) to beat standard 30s hardware NAT/firewall UDP timeouts
      - WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL=25s
      # MTU Fix (Critical): Prevent packet fragmentation from double-encryption
      - WIREGUARD_MTU=1280
      # Give WARP more time to establish the connection before internal restart (default is 6s which is too short)
      - HEALTH_VPN_DURATION_INITIAL=30s
      # Use plaintext DNS resolvers instead of DNS-over-TLS (DoT) to prevent startup failures on networks where port 853 is blocked
      - DNS_UPSTREAM_RESOLVER_TYPE=plain
      - DNS_UPSTREAM_PLAIN_ADDRESSES=1.1.1.1:53,8.8.8.8:53
      # Allow local network traffic to bypass WARP so Tailscale can establish fast, direct peer-to-peer connections (no DERP relays) on the LAN
      - FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8
      # Built-in DNS blocking for ads, malicious domains, and tracking
      # NOTE: This is the massive "Hagezi-style" built-in blocklist.
      # You can turn these to 'on' for maximum security if you have spare RAM
      # or are running this on a cloud exit node with adequate autonomous resources (requires approx 1-2GB of RAM).
      # Disabled intentionally: AdGuard Home provides the DNS filtering layer instead.
      # If AdGuard fails, there is no fallback filtering, but the healthcheck
      # dependency chain ensures Tailscale won't come up if AdGuard is down.
      - BLOCK_MALICIOUS=off
      - BLOCK_SURVEILLANCE=off
      - BLOCK_ADS=off
    dns:
      - 1.1.1.1
    sysctls:
      - net.ipv4.ip_forward=1
    volumes:
      - ./post-rules.txt:/iptables/post-rules.txt:ro
    restart: unless-stopped
    healthcheck:
      test: ["CMD-SHELL", "/gluetun-entrypoint healthcheck"]
      interval: 2s
      timeout: 5s
      retries: 15
      start_period: 60s
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

  tailscale:
    image: tailscale/tailscale:v1.98.4
    container_name: tailscale
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
      adguardhome:
        condition: service_healthy
    cap_add:
      - NET_ADMIN
      - NET_RAW
      - SYS_MODULE
    sysctls:
      - net.ipv4.ip_forward=1
      - net.ipv6.conf.all.forwarding=1
    environment:
      - TZ=UTC
      - TS_EXTRA_ARGS=--advertise-exit-node
      - PORT=41642
      - TS_USERSPACE=false
      - TS_AUTHKEY=${TS_AUTHKEY}
      - TS_STATE_DIR=/var/lib/tailscale
      # (Note: For headless/cloud deployments, you can bypass this footgun by
      # using Tailscale Ephemeral Keys which auto-clean themselves).
      - TS_AUTH_ONCE=true
      # MTU Optimization: Force Tailscale packets to be smaller than WARP's 1280 MTU
      # to prevent WireGuard-in-WireGuard fragmentation, completely eliminating bufferbloat
      - TS_DEBUG_MTU=1200
      # Enable Tailscale's built-in SOCKS5 proxy to route out of the WARP tunnel
      - TS_SOCKS5_SERVER=:1080
    devices:
      - /dev/net/tun:/dev/net/tun
    volumes:
      - tailscale_state:/var/lib/tailscale
    restart: unless-stopped
```

```
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

  routing-fix:
    image: nullexit-routing-fix:v1.0.0
    build:
      context: ./scripts
      dockerfile: Dockerfile.routing-fix
    container_name: routing-fix
    network_mode: service:warp
    # NOTE: Deliberately NOT sharing tailscale's PID namespace  the shared
    # namespace caused routing-fix to be killed whenever tailscale restarted
    # (e.g. during gluetun health-check flaps), which dropped all routing
    # table changes and left traffic leaking through eth0 instead of tun0.
    cap_add:
      - NET_ADMIN
      - SYSLOG
    depends_on:
      warp:
        condition: service_healthy
      # rule-compiler dependency removed: it is now a hash-gated one-off run by
      # toggle.sh before `up` (15.8.8). This service boots on the existing
      # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
    environment:
      - TZ=UTC
      - WARP_ENDPOINT=${WARP_ENDPOINT_1:-162.159.192.1}
      - BLOCKED_COUNTRIES=${BLOCKED_COUNTRIES}
      - TAILSCALE_ALLOW_LAN_P2P=${TAILSCALE_ALLOW_LAN_P2P:-false}
    volumes:
      - ./scripts/routing-fix.sh:/routing-fix.sh:ro
      # SECURITY: ./adguard/work/ is bind-mounted into this container.
      # Never write to it from the host or another container. The
      # single allowed writer is `rule-compiler` container  manual
      # edits corrupt the AdGuard cache + BoltDB session store.
      - ./adguard/work/userfilters:/userfilters:ro
      - ./adguard/work:/adguard_work:ro
      - .:/app
    restart: unless-stopped
    stop_grace_period: 1s
    logging:
      driver: "json-file"
      options:
        max-size: "1m"
        max-file: "1"

  adguardhome:
    # WARNING: AdGuard is pre-configured with hardcoded credentials (admin/nullexit).
    # This is fine for a local Mac running behind a NAT, but if you are deploying
    # this on a public cloud VPS with port 80/3000 exposed, change this immediately!
    image: adguard/adguardhome:v0.107.77
    container_name: adguardhome
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
      # rule-compiler dependency removed: it is now a hash-gated one-off run by
      # toggle.sh before `up` (15.8.8). This service boots on the existing
      # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
    environment:
      - TZ=UTC
    healthcheck:
      test: ["CMD-SHELL", "nc -z 127.0.0.1 5335 || exit 1"]
      interval: 2s
      timeout: 5s
      retries: 10
    volumes:
      # SECURITY: ./adguard/work/ is bind-mounted into this container.
      # Never write to it from the host or another container. The
      # single allowed writer is `rule-compiler` container  manual
      # edits corrupt the AdGuard cache + BoltDB session store.
      - ./adguard/work:/opt/adguardhome/work
      - ./adguard/conf:/opt/adguardhome/conf
    restart: unless-stopped
    stop_grace_period: 2s
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
```

```
        max-file: "3"

  rule-compiler:
    image: nullexit-rule-compiler:v1.0.0
    build:
      context: ./scripts
      dockerfile: Dockerfile.rule-compiler
    container_name: rule-compiler
    # `compile` profile: excluded from `docker compose up`. toggle.sh runs it
    # on-demand via `docker compose run --build --rm rule-compiler`, hash-gated on
    # black_list.txt/white_list.txt, so ~340k rules are re-deduped only when a list
    # actually changes instead of every toggle (~28s saved on the common path, 15.8.8).
    profiles: ["compile"]
    volumes:
      - .:/app

  tor:
    build: ./docker/tor
    image: nullexit-tor:v1.0.0
    container_name: tor
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
    environment:
      - TZ=UTC
      - TOR_USE_BRIDGES=${TOR_USE_BRIDGES:-false}
      - TOR_BRIDGE_LINES_FILE=${TOR_BRIDGE_LINES_FILE:-./tor-bridges.txt}
      - TOR_PASSWORD=${ADGUARD_PASSWORD:-nullexit}
      - TOR_EXCLUDE_EXIT_NODES=${TOR_EXCLUDE_EXIT_NODES:-{us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il}}
    working_dir: /nullexit
    volumes:
      - ./:/nullexit:ro
      - ./output.log:/output.log:rw
    restart: unless-stopped
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

volumes:
  tailscale_state:

networks:
  default:
    ipam:
      config:
        - subnet: 10.200.1.0/24
```

11 scripts/common.sh

```
#!/bin/bash
# common.sh  shared bash utilities for nullexit scripts

#  Formatting
RED='\033[0;31m'
GREEN='\033[0;32m'
YELLOW='\033[1;33m'
BLUE='\033[0;34m'
BOLD='\033[1m'
NC='\033[0m'

#  Constants & Environment
export MARKER_FILE="/tmp/nullexit-gateway-active.marker"
export PID_CAFFEINATE="/tmp/nullexit-caffeinate.pid"
export PID_DNS_WATCHER="/tmp/nullexit-dns-watcher.pid"
export DEFAULT_WARP_ENDPOINT_1="162.159.192.1"
export DEFAULT_WARP_ENDPOINT_2="162.159.193.1"

setup_standard_path() {
  export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
}

step() { echo -e "\n[$(date '+%Y-%m-%d %H:%M:%S')] ${BLUE}${BOLD} $*${NC}"; }
ok()   { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${GREEN} $*${NC}"; }
```

```
warn() { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${YELLOW} $*${NC}"; }
fail() { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${RED} $*${NC}"; }
die()  { echo -e "\n[$(date '+%Y-%m-%d %H:%M:%S')]   ${RED} $*${NC}\n"; exit 1; }

#  Helper Functions

# Append a timestamped line to output.log (best-effort; never fails the caller).
# Used to record lifecycle breadcrumbs (EXEC/EXIT markers, phase changes) so a
# START/STOP that is killed or window-closed still leaves a retraceable trail.
log_line() {
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] $*" >> "${LOG_FILE:-output.log}" 2>/dev/null || true
}

# Run a lifecycle command with FULL observability: log the command, stream its
# combined stdout+stderr to the terminal AND append it to output.log via tee,
# then log the exit code explicitly. Returns the command's real exit code (not
# tee's). This closes the blind spot where `docker compose up` / `colima start`
# printed only to the terminal  so when they fail (or the run is interrupted),
# output.log shows exactly what the last command was and whether it succeeded.
#
# Usage:  run_logged docker compose up -d --build
run_logged() {
  local logf="${LOG_FILE:-output.log}"
  log_line "EXEC: $*"
  # PIPESTATUS[0] is the command's exit code; tee (PIPESTATUS[1]) is ignored.
  "$@" 2>&1 | tee -a "$logf"
  local rc=${PIPESTATUS[0]}
  log_line "EXIT $rc: $*"
  return "$rc"
}

# Pure-bash timeout (no dependency on GNU coreutils' timeout)
# Runs a command with a safety cutoff so a wedged daemon can't hang the script.
run_with_timeout() {
  local timeout_sec="$1"
  shift
  if [ $# -eq 0 ]; then return 1; fi

  "$@" &
  local cmd_pid=$!
  CURRENT_BG_PID="$cmd_pid"

  (
    sleep "$timeout_sec"
    if kill -0 "$cmd_pid" 2>> output.log; then
      echo -e "\n[Timeout] Command '$*' exceeded $timeout_sec seconds. Terminating..." >&2
      kill -15 "$cmd_pid" 2>> output.log || true
      sleep 2
      if kill -0 "$cmd_pid" 2>> output.log; then
        kill -9 "$cmd_pid" 2>> output.log || true
      fi
    fi
  ) &
  local watcher_pid=$!

  # Disable set -e temporarily to capture exit status of wait
  set +e
  wait "$cmd_pid" 2>> output.log
  local exit_status=$?
  set -e

  # Kill the watcher since the command finished
  kill "$watcher_pid" 2>> output.log && wait "$watcher_pid" 2>> output.log || true

  CURRENT_BG_PID=""
  return "$exit_status"
}

# Start Colima and return the moment Docker is actually reachable, instead of
# blocking on colima 0.10.3's post-provision hang. Observed on macOS/vz: the VM
# reaches READY and creates the `colima` docker context in ~10-15s, then
# `colima start` frequently fails to return for ~110s (until a watchdog kills it)
# even though Docker is fully usable the entire time. We poll the colima docker
# context and proceed as soon as it answers, then reap the (usually hung) starter.
# The VM keeps running and `colima status` stays accurate. $@ = colima start args.
# Returns 0 once Docker responds, 1 if it never does within the cap.
colima_start_until_ready() {
  local max_wait=90 waited=0 ready=false start_pid
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] EXEC(bg): colima start $*" >> output.log
  colima start "$@" >> output.log 2>&1 &
  start_pid=$!
```

```
  while [ "$waited" -lt "$max_wait" ]; do
    if run_with_timeout 5 docker --context colima info >/dev/null 2>&1; then
      ready=true
      break
    fi
    # If colima start exited on its own (genuine finish or hard failure), stop waiting.
    if ! kill -0 "$start_pid" 2>/dev/null; then
      run_with_timeout 5 docker --context colima info >/dev/null 2>&1 && ready=true
      break
    fi
    sleep 3
    waited=$((waited + 3))
  done
  # Reap the starter (harmless if it already exited); this does NOT stop the VM.
  kill "$start_pid" 2>/dev/null || true
  wait "$start_pid" 2>/dev/null || true
  if [ "$ready" = "true" ]; then
    docker context use colima >/dev/null 2>&1 || true
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker ready after ~${waited}s (starter reaped)." >> output.log
    return 0
  fi
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker NOT ready after ${max_wait}s." >> output.log
  return 1
}

# Hash the files that feed the long-running locally-built images that `docker
# compose up` starts (routing-fix, tor). toggle.sh uses this to skip
# `docker compose --build` when nothing changed: BuildKit re-evaluating those
# images on the 600MB Colima VM costs ~30s per toggle even when every layer is
# CACHED (15.8.8). Includes the Dockerfiles, everything they COPY (routing-fix.sh,
# tor/entrypoint.sh, the logger/ Go source), and docker-compose.yml. The
# rule-compiler image is NOT here  it is a `compile`-profiled one-off, built+run
# by the gated compile step (15.8.8). Prints a sha256 hex digest, or empty.
compute_build_inputs_hash() {
  local root="${SCRIPT_DIR:-.}" f
  {
    for f in \
      "$root/scripts/Dockerfile.routing-fix" \
      "$root/scripts/routing-fix.sh" \
      "$root/docker/tor/Dockerfile" \
      "$root/docker/tor/entrypoint.sh" \
      "$root/docker-compose.yml"; do
      printf '::%s::' "$f"; cat "$f" 2>/dev/null
    done
    find "$root/scripts/logger" -type f 2>/dev/null \
      | LC_ALL=C sort | while IFS= read -r f; do printf '::%s::' "$f"; cat "$f" 2>/dev/null; done
  } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
}

# Hash the ad-block lists the user controls: black_list.txt (custom blocks) and
# white_list.txt (the allow-list that governs what DNS resolves through). toggle.sh
# gates the ~28s rule recompile on this  editing either list changes the digest
# and triggers a one-off recompile of compiled_rules.txt + ip_blocklist.ipset;
# an unchanged toggle skips it entirely (15.8.8). Prints a sha256 hex digest.
compute_rules_hash() {
  local root="${SCRIPT_DIR:-.}"
  { printf '::black::'; cat "$root/black_list.txt" 2>/dev/null
    printf '::white::'; cat "$root/white_list.txt" 2>/dev/null
  } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
}

# Check if the gateway is active (either containers are running, or host DNS is hijacked)
is_gateway_active() {
  # 1. Check if containers are running (suppress stderr to avoid errors if docker is down)
  if run_with_timeout 15 docker compose ps --status running 2>/dev/null | grep -q 'warp'; then
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (warp container is running)" >> "$LOG_FILE"
    return 0
  fi

  # 2. Check if host DNS was hijacked (not 1.1.1.1 and not default/empty)
  local current_dns=""
  local active_svc=""
  if [[ "$OSTYPE" == "darwin"* ]]; then
    active_svc=$(get_active_service)
    current_dns=$(networksetup -getdnsservers "$active_svc" 2>> output.log || true)
  else
    current_dns=$(resolvectl dns 2>/dev/null | awk '/Global|Link/ {for(i=4;i<=NF;i++) print $i}' | head -n1)
  fi

  if [[ -n "$current_dns" && "$current_dns" != "1.1.1.1" && ! "$current_dns" =~ "There aren't any DNS Servers"
     ]]; then
```

```
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (DNS was hijacked: $current_dns)" >> "
     $LOG_FILE"
    return 0
  fi

  # 3. Check if SOCKS proxy is enabled (macOS only)
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local socks_proxy
    socks_proxy=$(networksetup -getsocksfirewallproxy "$active_svc" 2>> output.log || true)
    if echo "$socks_proxy" | grep -q "Enabled: Yes"; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (SOCKS proxy enabled)" >> "$LOG_FILE"
      return 0
    fi
  fi

  echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: FALSE (Containers down, DNS clean, SOCKS disabled)"
     >> "$LOG_FILE"
  return 1
}

# Decide whether the gateway SHOULD be running, from persisted *intent* rather
# than a live health probe. This is the correct signal for toggle.sh's
# STOP-vs-START decision: a disable should tear down whatever the last START
# left behind  it should NOT depend on whether Docker/Colima happens to be
# reachable right now.
#
# Why this exists: using is_gateway_active() for the decision means every toggle
# (including *disable*) runs `docker compose ps` first. When Colima is down the
# socket is absent, that call blocks for the full run_with_timeout window, and
# under `set -e` + the ERR trap a non-zero result at the `if  then  else  fi`
# boundary could abort the script BEFORE the START body runs  so the very path
# that would boot Colima never executed (observed as `EXIT: code=1 phase=start`).
#
# Signal source: MARKER_FILE is written at the end of a successful START
# (write_gateway_active_marker) and cleared at the top of STOP and in
# cleanup_handler. Reading it is instant and cannot hang or abort.
#
# ECHOES "running" or "stopped" to stdout (and ALWAYS returns 0). Callers read
# the string:  [ "$(gateway_state)" = "running" ] &&
#
# CRITICAL  why this echoes instead of using a 0/1 return code: on macOS
# /bin/bash 3.2, with `set -e` + an `ERR` trap + `set -x` (DEBUG_TRACE=true) all
# active, a function that does `return 1` triggers a SPURIOUS `set -e` abort in
# the caller even when the call is guarded (`if f; then`, or even `f || rc=$?`).
# The `return N` from the function, traced by `set -x`, trips the ERR trap. This
# is a genuine bash-3.2 interaction (see devref 15.11.8). Echoing a result and
# always returning 0 sidesteps `return`/`set -e`/`set -x` entirely  the function
# can never contribute a non-zero status, so no caller can be aborted by it.
gateway_state() {
  # Primary: persisted intent. Marker present => a START completed and no STOP
  # has cleared it yet.
  #
  # NOTE: toggle.sh runs a "defensive stale-marker clear" near the top of every
  # invocation (before this decision) to stop the wake-watcher firing against a
  # crashed gateway. That would wipe the marker before we read it, so toggle.sh
  # captures the marker's existence into GATEWAY_MARKER_AT_STARTUP *before* that
  # clear and we honor it here. If the variable is unset (other callers, or the
  # --restart pre-check which runs before the defensive clear), fall back to the
  # live marker file.
  if [ -n "${GATEWAY_MARKER_AT_STARTUP:-}" ]; then
    if [ "$GATEWAY_MARKER_AT_STARTUP" = "yes" ]; then echo "running"; return 0; fi
  elif [ -f "$MARKER_FILE" ]; then
    echo "running"; return 0
  fi

  # Marker absent. Normally this means "stopped". But cover the edge case where
  # the marker was lost (e.g. /tmp cleared, or gateway started by an older build)
  # yet containers are genuinely up  reconcile via is_gateway_active, but ONLY
  # when Docker is actually reachable, so a dead-socket probe can never hang. On
  # macOS Docker runs in the Colima VM; if that socket is absent, the gateway
  # definitively is not running.
  local docker_up="no"
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if [ -S /var/run/docker.sock ] || [ -S "$HOME/.colima/default/docker.sock" ]; then docker_up="yes"; fi
  else
    if [ -S /var/run/docker.sock ]; then docker_up="yes"; fi
  fi

  if [ "$docker_up" = "yes" ] && is_gateway_active; then
    echo "running"; return 0
  fi
```

```
  echo "stopped"
  return 0
}

# Reads an environment variable from a file (default: .env), strips quotes and whitespace
read_env_var() {
  local var_name="$1"
  local env_file="${2:-.env}"
  grep -E "^${var_name}=" "$env_file" 2>/dev/null | cut -d'=' -f2- | tr -d "\"\\'" | sed -e 's/^[[:space:]]*//' -e 's/[[:space:]]*$//' || echo ""
}

# Load KILL_SWITCH variable from .env (defaults to false if not found/empty)
export KILL_SWITCH
KILL_SWITCH=$(read_env_var "KILL_SWITCH" 2>/dev/null | tr '[:upper:]' '[:lower:]')
if [ -z "$KILL_SWITCH" ]; then
  KILL_SWITCH="false"
fi

# Function to get the active network service name (e.g., "Wi-Fi" or "USB 10/100 LAN")
get_active_service() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local iface
    iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')

    # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
    if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
      for i in en0 en1 en2 en3; do
        if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "inet "; then
          iface="$i"
          break
        fi
      done
    fi

    # Default fallback if still empty or not enX
    if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
      iface="en0"
    fi

    # Map interface (e.g., en0) to service name (e.g., Wi-Fi)
    local service
    service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B 1 "Device: $iface" | head -n 1 | sed -E 's/^\([0-9\*]+\) //')

    if [ -n "$service" ]; then
      echo "$service"
    else
      echo "Wi-Fi"
    fi
  else
    # Get the interface used for default internet routing
    local iface
    iface=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
    if [ -z "$iface" ]; then
      echo ""
      return 1
    fi
    echo "$iface"
  fi
}

get_en0_service() {
  local service
  service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B1 "Device: en0" | head -1 | sed -E 's/^\([0-9\*]+\) //' || true)
  if [ -z "$service" ]; then
    echo "Wi-Fi"
  else
    echo "$service"
  fi
}

get_warp_endpoint_1() { read_env_var "WARP_ENDPOINT_1" ".env" | grep -Eo '^[0-9\.]+' || echo "$DEFAULT_WARP_ENDPOINT_1"; }
get_warp_endpoint_2() { read_env_var "WARP_ENDPOINT_2" ".env" | grep -Eo '^[0-9\.]+' || echo "$DEFAULT_WARP_ENDPOINT_2"; }

restart_tailscaled_daemon() {
  echo "  [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
```

```
  echo "  [Tailscale Recovery] Attempting to restart tailscaled service..."

  if run_with_timeout 15 brew services restart tailscale >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled (user service)."
  elif run_with_timeout 15 sudo -n brew services restart tailscale >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled (system service)."
  else
    echo "  [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
  fi
  sleep 3
}

disconnect_tailscale_host() {
  local ts_bin="${1:-tailscale}"
  if command -v "$ts_bin" >> output.log 2>&1; then
    echo "  Disconnecting host Tailscale from mesh..."

    # Check if actually connected first (with timeout  tailscaled could be wedged)
    local status_ok=false
    if run_with_timeout 5 "$ts_bin" status >> output.log 2>&1; then
      status_ok=true
    else
      local status_exit=$?
      # If it was a quick exit code 1 (disconnected), it is still reachable
      if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
        status_ok=true
      fi
    fi

    if [ "$status_ok" = "false" ]; then
      restart_tailscaled_daemon
    fi

    # Explicitly reset any exit-node so it doesn't linger.
    "$ts_bin" up >> output.log 2>&1 || true
    if run_with_timeout 10 "$ts_bin" set --ssh=true --accept-dns=false --exit-node= >> output.log 2>&1; then
      echo "  [] Exit-node preference cleared."
    else
      echo "  [!] tailscale up didn't respond (tailscaled may be wedged)"
    fi

    local ts_args=""
    if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
    if run_with_timeout 10 "$ts_bin" down $ts_args >> output.log 2>&1; then
      echo "  [] Tailscale disconnected."
    else
      echo "  [!] tailscale down didn't respond"
    fi
  else
    echo "  [!] tailscale CLI not found  skipping"
  fi
}

stop_pidfile_daemon() {
  local pid_file="$1"
  local label="$2"
  if [ -f "$pid_file" ]; then
    local pid
    pid=$(cat "$pid_file")
    if kill -0 "$pid" 2>/dev/null; then
      echo "  Stopping $label (PID $pid)..."
      sudo -n kill "$pid" 2>> output.log || kill "$pid" 2>> output.log || true
      sleep 0.5
      if kill -0 "$pid" 2>/dev/null; then
        sudo -n kill -9 "$pid" 2>> output.log || kill -9 "$pid" 2>> output.log || true
      fi
    fi
    rm -f "$pid_file"
  fi
}

disable_all_proxies() {
  local svc="$1"
  if [ -n "$svc" ]; then
    sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
    sudo -n networksetup -setwebproxystate "$svc" off 2>> output.log || true
    sudo -n networksetup -setsecurewebproxystate "$svc" off 2>> output.log || true
  fi
}

flush_dns_cache() {
```

```
  if [[ "$OSTYPE" == "darwin"* ]]; then
    sudo -n dscacheutil -flushcache 2>> output.log || true
    sudo -n killall -HUP mDNSResponder 2>> output.log || true
  else
    resolvectl flush-caches 2>> output.log || true
  fi
}

wait_for_dhcp_settle() {
  # Resolve physical interface (Wi-Fi or Ethernet)
  local physical_iface
  physical_iface=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
     getline; print $2; exit}')
  [ -z "$physical_iface" ] && physical_iface="en0"

  echo -n "Waiting for physical network interface ($physical_iface) DHCP lease to settle..."
  local settled=false
  for attempt in {1..60}; do
    PHYSICAL_GW=$(ipconfig getpacket "$physical_iface" 2>> output.log | awk -F'[{}]' '/router /{print $2}')
    local local_ip
    local_ip=$(ifconfig "$physical_iface" 2>/dev/null | awk '/inet /{print $2}')

    if [ -n "$PHYSICAL_GW" ] && [[ "$PHYSICAL_GW" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
       [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
       [[ "$local_ip" != "169.254."* ]]; then
      echo " settled (Interface IP: $local_ip, Router IP: $PHYSICAL_GW)."
      settled=true
      break
    fi

    if [ "$attempt" -gt 15 ] && [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] &&
     [[ "$local_ip" != "169.254."* ]]; then
      local fallback_gw
      fallback_gw=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
      if [ -n "$fallback_gw" ] && [[ "$fallback_gw" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]]; then
        PHYSICAL_GW="$fallback_gw"
      fi
      echo " static/settled detected (Interface IP: $local_ip, Gateway IP: ${PHYSICAL_GW:-unknown})."
      settled=true
      break
    fi
    echo -n "."
    sleep 0.5
  done
  if [ "$settled" = "false" ]; then
    echo " timed out or no active link. Proceeding anyway."
  fi
}

reset_sharing_services() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    sudo -n killall sharingd rapportd 2>> output.log || true
  fi
}

read_adguard_ip() {
  if [ -f ".gateway_ip" ]; then
    tr -d '\r' < ".gateway_ip" | awk 'NR==1{print $1;exit}'
  fi
}

is_kill_switch_enabled() {
  [ "$KILL_SWITCH" = "true" ]
}

add_warp_bypass_routes() {
  local target="${1:-}"
  local ep1
  ep1=$(get_warp_endpoint_1)
  local ep2
  ep2=$(get_warp_endpoint_2)

  if [[ "$OSTYPE" == "darwin"* ]]; then
    local routing_arg=""
    local msg_via=""
    if [ -n "$target" ]; then
      if [[ "$target" =~ ^en[0-9]+$ || "$target" == "bridge"* ]]; then
        routing_arg="-interface $target"
        msg_via="interface $target"
      else
        routing_arg="$target"
```

```
        msg_via="gateway $target"
      fi
    else
      local gateway_ip
      gateway_ip=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
      if [ -z "$gateway_ip" ]; then
        local iface
        iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
        if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
          iface="en0"
        fi
        routing_arg="-interface $iface"
        msg_via="interface $iface"
      else
        routing_arg="$gateway_ip"
        msg_via="gateway $gateway_ip"
      fi
    fi

    echo -e "\nAdding host bypass routes for Cloudflare WARP endpoints via $msg_via..."
    sudo -n route delete -host "$ep1" 2>/dev/null || true
    sudo -n route delete -host "$ep2" 2>/dev/null || true
    sudo -n route add -host "$ep1" $routing_arg >> output.log 2>&1 || true
    sudo -n route add -host "$ep2" $routing_arg >> output.log 2>&1 || true

    echo -e "Adding host bypass routes for Tailscale control plane via $msg_via..."
    sudo -n route delete -net 192.200.0.0/24 2>/dev/null || true
    sudo -n route add -net 192.200.0.0/24 $routing_arg >> output.log 2>&1 || true

    echo -e "Adding host bypass routes for Tailscale DERP relays via $msg_via..."
    local derp_ips
    derp_ips=$(curl -s --connect-timeout 5 https://login.tailscale.com/derpmap/default | grep -oE '"IPv4"[[:
     space:]]*:[[:space:]]*"[0-9.]+"' | cut -d'"' -f4 | grep -E '^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$' || true)
    if [ -n "$derp_ips" ]; then
      echo "$derp_ips" > /tmp/nullexit-derp-ips.txt
      for ip in $derp_ips; do
        sudo -n route delete -host "$ip" 2>/dev/null || true
        sudo -n route add -host "$ip" $routing_arg >> output.log 2>&1 || true
      done
    fi
  else
    local gateway_ip="${target}"
    if [ -z "$gateway_ip" ]; then
      gateway_ip=$(ip route show default 2>/dev/null | awk '/default/ {print $3}')
    fi
    if [ -n "$gateway_ip" ]; then
      sudo ip route add "$ep1" via "$gateway_ip" >> output.log 2>&1 || true
      sudo ip route add "$ep2" via "$gateway_ip" >> output.log 2>&1 || true
    fi
  fi
}

remove_warp_bypass_routes() {
  local ep1
  ep1=$(get_warp_endpoint_1)
  local ep2
  ep2=$(get_warp_endpoint_2)

  if [[ "$OSTYPE" == "darwin"* ]]; then
    echo -e "\nRemoving host bypass routes for Cloudflare WARP endpoints..."
    sudo -n route delete -host "$ep1" >> output.log 2>&1 || true
    sudo -n route delete -host "$ep2" >> output.log 2>&1 || true
    echo -e "Removing host bypass routes for Tailscale control plane..."
    sudo -n route delete -net 192.200.0.0/24 >> output.log 2>&1 || true
    if [ -f /tmp/nullexit-derp-ips.txt ]; then
      echo -e "Removing host bypass routes for Tailscale DERP relays..."
      while read -r ip; do
        sudo -n route delete -host "$ip" >> output.log 2>&1 || true
      done < /tmp/nullexit-derp-ips.txt
      rm -f /tmp/nullexit-derp-ips.txt
    fi
  else
    sudo ip route del "$ep1" >> output.log 2>&1 || true
    sudo ip route del "$ep2" >> output.log 2>&1 || true
  fi
}

#  FaceTime / real-time P2P split-route (opt-in, default OFF)
# Mitigation for the FaceTime Double-Tunnel bug (15.12.22). Adds more-specific
# host routes for Apple's FaceTime media/relay /16s via the PHYSICAL gateway, so
# longest-prefix match beats the exit-node 0.0.0.0/1 + 128.0.0.0/1 and ONLY those
```

```
# /16s leave direct  real-time media can then hole-punch from the real IP instead
# of dying behind WARP's symmetric NAT. Mirrors add_warp_bypass_routes.
#
#  PRIVACY: the direct /16s expose the real ISP IP (not WARP) for that traffic
# a deliberate, narrow hole in the Double-Tunnel promise, scoped to Apple infra
# (which already ties traffic to your Apple ID). Inert unless FACETIME_SPLIT_ROUTE
# =true. Needs BOTH a route AND a PF pass: 'block out on en0 all' would otherwise
# drop the packets, so we populate the <apple_direct> table pf.conf permits.
facetime_split_enabled() {
  [ "$(read_env_var FACETIME_SPLIT_ROUTE | tr '[:upper:]' '[:lower:]')" = "true" ]
}

facetime_direct_subnets() {
  local s; s=$(read_env_var FACETIME_DIRECT_SUBNETS)
  echo "${s:-17.249.0.0/16}"
}

add_apple_split_routes() {
  facetime_split_enabled || return 0
  [[ "$OSTYPE" == "darwin"* ]] || return 0
  local gw="${1:-}"
  if [ -z "$gw" ]; then
    # The physical default route survives alongside the exit-node /1 override,
    # so the real gateway is still the 'default' entry in the table.
    gw=$(netstat -rn -f inet 2>/dev/null | awk '$1=="default"{print $2; exit}')
  fi
  if [ -z "$gw" ] || [[ "$gw" =~ ^utun ]] || [[ "$gw" =~ ^link ]]; then
    echo "  [FaceTime split-route] could not resolve a physical gateway  skipping."
    return 0
  fi
  local subnets; subnets=$(facetime_direct_subnets)
  echo -e "\n[FaceTime split-route] routing Apple subnets DIRECT via $gw (real IP exposed for these): $subnets"
  for net in $subnets; do
    sudo -n route delete -net "$net" >> output.log 2>&1 || true
    sudo -n route add -net "$net" "$gw" >> output.log 2>&1 || true
    # Kill-switch pass: the <apple_direct> table is what pf.conf lets out on en0.
    sudo -n pfctl -a com.apple/nullexit -t apple_direct -T add "$net" >> output.log 2>&1 || true
  done
}

remove_apple_split_routes() {
  [[ "$OSTYPE" == "darwin"* ]] || return 0
  local subnets; subnets=$(facetime_direct_subnets)
  for net in $subnets; do
    sudo -n route delete -net "$net" >> output.log 2>&1 || true
  done
  # Flush the whole table so no direct Apple exception lingers after teardown.
  sudo -n pfctl -a com.apple/nullexit -t apple_direct -T flush >> output.log 2>&1 || true
}

bounce_wifi_interfaces() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local wifi_port
    wifi_port=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print
     $2}')
    if [ -n "$wifi_port" ]; then
      sudo -n networksetup -setairportpower "$wifi_port" off 2>> output.log || true
      sleep 2
      sudo -n networksetup -setairportpower "$wifi_port" on 2>> output.log || true
      echo "Wi-Fi bounced (interface $wifi_port)"
      sleep 3
    else
      echo "Could not detect Wi-Fi interface. Bouncing fallback interfaces..."
      set +e
      for iface in en0 en1 en2 en3 en4 en5; do
        if ifconfig "$iface" >> output.log 2>&1; then
          sudo -n ifconfig "$iface" down 2>> output.log || true
          sudo -n ifconfig "$iface" up 2>> output.log || true
        fi
      done
      set -e
      echo "Fallback interfaces bounced"
    fi
  fi
}

get_active_interface() {
  local iface
  iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')

  # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
```

```
  if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
    for i in en0 en1 en2 en3; do
      if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "
     inet "; then
        iface="$i"
        break
      fi
    done
  fi

  # Default fallback if still empty or not enX
  if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
    iface="en0"
  fi

  echo "$iface"
}

enable_killswitch() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if ! is_kill_switch_enabled; then
      return 0
    fi

    local ext_if
    ext_if=$(get_active_interface)
    local ep1
    ep1=$(get_warp_endpoint_1)
    local ep2
    ep2=$(get_warp_endpoint_2)
    local pf_conf_path
    pf_conf_path="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/pf.conf"

    # Keep the host's scrub max-mss in sync with GATEWAY_MSS (same variable that drives
    # post-rules.txt's --set-mss for the container). The host egresses through the same
    # double-tunnel (Tailscale+WARP) as forwarded phone traffic, so it needs the same
    # empirically-safe clamp  a stale/higher value here reproduces the 15.3.2 stalling
    # bug for the host's own connections (see devref.md).
    local gateway_mss
    gateway_mss=$(read_env_var GATEWAY_MSS)
    if [[ "$gateway_mss" =~ ^[0-9]+$ ]]; then
      sed -i '' "s/max-mss [0-9]*/max-mss ${gateway_mss}/g" "$pf_conf_path" 2>/dev/null || true
    fi

    echo -e "\nEnabling macOS Packet Filter (PF) Kill-Switch on $ext_if..."

    # Enable PF and capture the reference token
    local token_out
    token_out=$(sudo -n pfctl -E 2>> output.log || true)
    local token
    token=$(echo "$token_out" | awk '/Token/{print $NF}')
    if [ -n "$token" ]; then
      echo "$token" > /tmp/nullexit-pf-token.txt
      echo "  [] PF enabled with token $token."
    else
      echo "  [!] WARNING: Could not capture PF reference token (may already be enabled or sudo failed)."
    fi

    # Load the rules into the anchor
    if sudo -n pfctl -D "ext_if=$ext_if" -a com.apple/nullexit -f "$pf_conf_path" >> output.log 2>&1; then
      echo "  [] Ruleset loaded into anchor com.apple/nullexit."
    else
      echo "  [!] ERROR: Failed to load ruleset into anchor."
    fi

    # Populate tables in the anchor
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T flush >> output.log 2>&1 || true
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep1" >> output.log 2>&1 || true
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep2" >> output.log 2>&1 || true

    if [ -f /tmp/nullexit-derp-ips.txt ]; then
      local derp_ips
      derp_ips=$(tr '\n' ' ' < /tmp/nullexit-derp-ips.txt)
      if [ -n "$derp_ips" ]; then
        sudo -n pfctl -a com.apple/nullexit -t derp_relays -T flush >> output.log 2>&1 || true
        sudo -n pfctl -a com.apple/nullexit -t derp_relays -T add $derp_ips >> output.log 2>&1 || true
      fi
    fi
    echo "  [] Kill-switch tables populated."
  fi
}
```

```
740
741 disable_killswitch() {
742   if [[ "$OSTYPE" == "darwin"* ]]; then
743     if ! is_kill_switch_enabled; then
744       return 0
745     fi
746     echo -e "\nDisabling macOS PF Kill-Switch..."
747
748     # Flush rules from anchor com.apple/nullexit
749     sudo -n pfctl -a com.apple/nullexit -F all >> output.log 2>&1 || true
750
751     # Release the enable reference if token is present
752     if [ -f /tmp/nullexit-pf-token.txt ]; then
753       local token
754       token=$(cat /tmp/nullexit-pf-token.txt)
755       if [ -n "$token" ]; then
756         sudo -n pfctl -X "$token" >> output.log 2>&1 || true
757         echo "  [] Released PF enable reference token $token."
758       fi
759       rm -f /tmp/nullexit-pf-token.txt
760     else
761       echo "  [] Rules flushed from anchor com.apple/nullexit."
762     fi
763   fi
764 }
765
766
767 write_host_ips() {
768   # Gather all active IPv4 addresses on the host and write to .host_ips
769   # routing-fix.sh will read this file to explicitly whitelist them for direct Tailscale P2P
770   local repo_dir
771   repo_dir="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"
772
773   if [[ "$OSTYPE" == "darwin"* ]]; then
774     ifconfig | awk '/inet /{print $2}' | grep -v '127.0.0.1' > "$repo_dir/.host_ips" 2>/dev/null || true
775   else
776     ip -4 addr show 2>/dev/null | awk '/inet /{print $2}' | cut -d/ -f1 | grep -v '127.0.0.1' > "$repo_dir/.
      host_ips" 2>/dev/null || true
777   fi
778 }
779
780 start_dns_watcher() {
781   if [[ "$OSTYPE" == "darwin"* ]]; then
782     local target_ip=$1
783     stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
784     echo "  Starting background DNS Watcher for seamless Wi-Fi roaming..."
785     nohup bash -c "
786       source \"$SCRIPT_DIR/scripts/common.sh\"
787       trap 'exit 0' SIGTERM SIGINT SIGHUP
788       while true; do
789         ACTIVE_IF=\$(get_active_service)
790         if [ -n \"\$ACTIVE_IF\" ]; then
791           CURRENT_DNS=\$(networksetup -getdnsservers \"\$ACTIVE_IF\" 2>/dev/null)
792           if [ \"\$CURRENT_DNS\" != \"$target_ip\" ]; then
793             networksetup -setdnsservers \"\$ACTIVE_IF\" \"$target_ip\" >/dev/null 2>&1
794           fi
795         fi
796         sleep 30
797       done
798     " >> output.log 2>&1 &
799     echo $! > "$DNS_WATCHER_PID_FILE"
800   else
801     local target_ip=$1
802     stop_dns_watcher
803     echo "  Starting background DNS Watcher for seamless Wi-Fi roaming..."
804     nohup bash -c "
805       trap 'exit 0' SIGTERM SIGINT SIGHUP
806       while true; do
807         if [ -n \"\$ACTIVE_SERVICE\" ]; then
808           CURRENT_DNS=\$(resolvectl dns \"\$ACTIVE_SERVICE\" 2>/dev/null | grep -oE
    '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head -1)
809           if [ \"\$CURRENT_DNS\" != \"$target_ip\" ]; then
810             resolvectl dns \"\$ACTIVE_SERVICE\" \"$target_ip\" >/dev/null 2>&1 || true
811           fi
812         fi
813         sleep 30
814       done
815     " >> output.log 2>&1 &
816     echo $! > "$DNS_WATCHER_PID_FILE"
817   fi
818 }
```

```
819
820 #  SOCKS5 Proxy (WARP Tunnel)
821 # When Tailscale's exit node can't route through WARP (due to userspace
822 # forwarding), we use a SOCKS5 proxy running inside the warp container's
823 # network namespace. Connections created by this proxy go through the
824 # kernel routing table (table 200 -> tun0 -> WARP), unlike Tailscale's
825 # userspace exit node which bypasses it.
826 #
827 # The SOCKS5 proxy is exposed on localhost:${SOCKS_PROXY_PORT} via Docker port mapping.
828 # `networksetup -setsocksfirewallproxy` on macOS routes system TCP traffic
829 # through the proxy, which then goes through WARP.
830 #
831 # IMPORTANT: CLI tools like curl do NOT respect macOS system proxy settings.
832 # They must use --socks5-hostname or the ALL_PROXY env variable explicitly.
833
834 enable_socks_proxy() {
835   if [[ "$OSTYPE" == "darwin"* ]]; then
836     local svc="$1"
837     if [ -z "$svc" ]; then
838       svc="$ACTIVE_SERVICE"
839     fi
840
841     echo -n "  Enabling system-wide SOCKS5 proxy on $svc (localhost:$SOCKS_PROXY_PORT)... "
842     sudo -n networksetup -setsocksfirewallproxy "$svc" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.log || true
843     sudo -n networksetup -setsocksfirewallproxystate "$svc" on 2>> output.log || true
844
845     # Also set on en0 service for macOS resolver consistency
846     if [ "$svc" != "$EN0_SERVICE" ]; then
847       sudo -n networksetup -setsocksfirewallproxy "$EN0_SERVICE" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.log ||
      true
848       sudo -n networksetup -setsocksfirewallproxystate "$EN0_SERVICE" on 2>> output.log || true
849     fi
850
851     # IMPORTANT: Exclude local LAN and mDNS traffic from the proxy so P2P works natively
852     # without source-IP masking from the Colima VM bridge.
853     sudo -n networksetup -setproxybypassdomains "$svc" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
      172.16.0.0/12 "*.local" 2>> output.log || true
854     if [ "$svc" != "$EN0_SERVICE" ]; then
855       sudo -n networksetup -setproxybypassdomains "$EN0_SERVICE" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
      172.16.0.0/12 "*.local" 2>> output.log || true
856     fi
857
858     echo "done."
859     echo "  All TCP traffic now routed through gateway -> WARP tunnel -> internet."
860     echo "  (CLI tools: export ALL_PROXY=socks5://127.0.0.1:$SOCKS_PROXY_PORT)"
861   else
862     local svc="$1"
863     echo "  (Linux global SOCKS proxy configuration is desktop-environment specific, skipping.)"
864     echo "  SOCKS5 proxy is active and listening on 127.0.0.1:$SOCKS_PROXY_PORT"
865   fi
866 }
867
868 disable_socks_proxy() {
869   if [[ "$OSTYPE" == "darwin"* ]]; then
870     for svc in "$ACTIVE_SERVICE" "$EN0_SERVICE"; do
871       sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
872     done
873     echo "  SOCKS5 proxy disabled."
874   else
875     local svc="$1"
876     echo "  (Linux global SOCKS proxy configuration skipped.)"
877   fi
878 }
879
880 # Helper to restore host DNS to a clean state (used on script start, on failure,
881 # and after a successful STOP). Without this, a successful toggle-off leaves the
882 # host's resolver pinned to a now-dead gateway IP  every lookup stalls for macOS's
883 # DNS timeout (~5s) before falling through to whatever next server is configured.
884 # With it, we always leave the host in `1.1.1.1 / no search domain`.
885 #
886 # We set DNS on BOTH the active service AND the en0 service because macOS's scutil DNS
887 # resolver is commonly scoped to en0 (Wi-Fi). A per-service change on e.g.
888 # "USB 10/100 LAN" is ignored by the system resolver unless the en0 service is also updated.
889 reset_dns() {
890   if [[ "$OSTYPE" == "darwin"* ]]; then
891     if [ -n "$ACTIVE_SERVICE" ]; then
892       for dns_svc in "$ACTIVE_SERVICE" "$EN0_SERVICE"; do
893         networksetup -setsearchdomains "$dns_svc" "Empty" 2>> output.log || true
894         networksetup -setdnsservers "$dns_svc" 1.1.1.1 2>> output.log || true
895       done
896     fi
```

```
  else
    if [ -n "$ACTIVE_SERVICE" ]; then
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      resolvectl revert "$ACTIVE_SERVICE" 2>/dev/null || true
    fi
  fi
}

# Verify basic direct-internet connectivity after a recovery/teardown: DNS via
# 1.1.1.1, then HTTP (curl) or ping to 1.1.1.1. Sets the global INTERNET_OK to
# "true" on any successful check (callers read $INTERNET_OK in their summary).
# Always returns 0 so a bare call can't trip `set -e`. Used by recover.sh's macOS
# and Linux recovery paths (extracted from a formerly-duplicated block).
verify_internet_connectivity() {
  step "Verifying internet connectivity"

  # Wait a moment for the network to settle
  sleep 2

  INTERNET_OK=false

  # Try DNS resolution first
  if host -W 3 google.com 1.1.1.1 >> output.log 2>&1; then
    ok "DNS resolution works (google.com via 1.1.1.1)"
    INTERNET_OK=true
  elif nslookup google.com 1.1.1.1 >> output.log 2>&1; then
    ok "DNS resolution works (nslookup google.com via 1.1.1.1)"
    INTERNET_OK=true
  else
    warn "DNS resolution check failed  will still try ping/curl"
  fi

  # Try actual HTTP connectivity
  if command -v curl >> output.log 2>&1; then
    if curl -sf --max-time 5 https://1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via HTTP"
      INTERNET_OK=true
    elif curl -sf --max-time 5 http://1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via HTTP (plain)"
      INTERNET_OK=true
    else
      warn "HTTP check failed"
    fi
  elif command -v ping >> output.log 2>&1; then
    if ping -c 1 -W 3 1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via ping"
      INTERNET_OK=true
    else
      warn "Ping check failed"
    fi
  fi

  return 0
}

#  Local DNS Proxy (Python)
# When the tailnet data plane cannot establish, we use a Python DNS proxy.
# It listens on UDP:53, receives DNS queries, forwards them over TCP to
# AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (Docker port mapping), and returns the response.
# The Python script handles the DNS-over-TCP wire format (2-byte length prefix)
# correctly  unlike socat which just forwards raw bytes.
#
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_PID=""

# Kill any leftover DNS proxy process
stop_local_dns_proxy() {
  if [ -n "$DNS_PROXY_PID" ]; then
    kill "$DNS_PROXY_PID" 2>/dev/null || true
    wait "$DNS_PROXY_PID" 2>/dev/null || true
    DNS_PROXY_PID=""
  fi
  # Clean up any stale Python DNS proxy processes
  sudo -n pkill -f "dns-proxy.py" 2>/dev/null || true
  rm -f /tmp/nullexit-dns-proxy.conf 2>/dev/null || true
}

# Start local DNS proxy (Python)
start_local_dns_proxy() {
  if [ -z "$DNS_PROXY_BIN" ]; then
    echo "  Python not found. Python3 is built into macOS  this shouldn't happen."
```

```
    return 1
  fi

  if [ ! -f "$DNS_PROXY_SCRIPT" ]; then
    echo "  DNS proxy script not found at $DNS_PROXY_SCRIPT"
    return 1
  fi

  # Kill any leftover process first
  stop_local_dns_proxy

  echo -n "  Starting local DNS proxy via Python (UDP:53  TCP:${DNS_PROXY_PORT})... "
  # The shipped sudoers NOPASSWD rule only covers the literal `python3
  # dns-proxy.py` command  no bash -c wrapper, no env prefix  so sudo -n's
  # env_reset strips TARGET_PORT before it ever reaches the script. Drop it in
  # a file dns-proxy.py reads itself instead (see its own comment for why).
  echo "TARGET_PORT=$DNS_PROXY_PORT" > /tmp/nullexit-dns-proxy.conf
  chmod 600 /tmp/nullexit-dns-proxy.conf 2>/dev/null || true
  sudo -n "$DNS_PROXY_BIN" "$DNS_PROXY_SCRIPT" &
  DNS_PROXY_PID=$!
  disown "$DNS_PROXY_PID" 2>/dev/null || true

  # Give it a moment to bind
  sleep 0.5

  if kill -0 "$DNS_PROXY_PID" 2>/dev/null; then
    echo "started (PID $DNS_PROXY_PID)."

    # Hijack host DNS to localhost
    echo -n "  Hijacking host DNS to 127.0.0.1 for ad-blocking... "
    if [[ "$OSTYPE" == "darwin"* ]]; then
      networksetup -setsearchdomains "$ACTIVE_SERVICE" "Empty" 2>/dev/null || true
      networksetup -setsearchdomains "$ENO_SERVICE" "Empty" 2>/dev/null || true
      if ! networksetup -setdnsservers "$ACTIVE_SERVICE" 127.0.0.1; then
        echo "FAILED (networksetup error  check permissions)."
        stop_local_dns_proxy
        return 1
      fi
      networksetup -setdnsservers "$ENO_SERVICE" 127.0.0.1 2>/dev/null || true
      sudo -n dscacheutil -flushcache 2>/dev/null || true
      sudo -n killall -HUP mDNSResponder 2>/dev/null || true
    else
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" 127.0.0.1; then
        echo "FAILED (resolvectl error  check permissions)."
        stop_local_dns_proxy
        return 1
      fi
      resolvectl dns "$ACTIVE_SERVICE" 127.0.0.1 2>/dev/null || true
      sudo -n resolvectl flush-caches 2>/dev/null || true
    fi

    # Verify DNS actually works through the proxy
    echo -n "  Verifying DNS resolution... "
    if dig @127.0.0.1 google.com +short +timeout=5 &>/dev/null; then
      echo "ok."
      return 0
    else
      echo "FAILED (DNS queries not reaching AdGuard)."
      echo "  Restoring DNS to 1.1.1.1..."
      reset_dns
      stop_local_dns_proxy
      return 1
    fi
  else
    echo "FAILED (could not bind port 53  is it in use?)."
    DNS_PROXY_PID=""
    return 1
  fi
}

# Helper to FORCIBLY hijack host DNS to a single server (the gateway's AdGuard IP)
# and VERIFY via `networksetup -getdnsservers` that the only name server is exactly
# $1. Returns 0 iff the live state matches; non-zero if it can't be confirmed.
#
# This deliberately does NOT append a 1.1.1.1 fallback: macOS's resolver queries
# the list in order and falls back to the next entry on timeout, so on a
# misbehaving AdGuard (filter compile crash, OOM kill, etc.) the resolver still
# leaks to 1.1.1.1 and silently bypasses ad-blocking. With no fallback, a broken
# gateway manifests as a *visible* DNS outage that prompts the user to
```

```
# investigate  which is the correct failure mode.
force_dns_to_gateway() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local target_ip="$1"
    if [ -z "$target_ip" ] || [ -z "$ACTIVE_SERVICE" ]; then
      echo "  [force_dns] ERROR: target_ip or ACTIVE_SERVICE is empty" >&2
      return 1
    fi

    local attempt
    for attempt in 1 2 3; do
      echo -n "  [force_dns] Attempt $attempt/3: setting DNS to $target_ip on \"$ACTIVE_SERVICE\" + \"
     $EN0_SERVICE\"... "

      # Apply to BOTH the active service AND the en0 service  the scutil DNS resolver
      # is scoped to en0 (Wi-Fi) and ignores per-service changes that don't include it.
      # Fail fast on ACTIVE_SERVICE; best-effort on EN0_SERVICE (it may not exist).
      networksetup -setsearchdomains "$ACTIVE_SERVICE" "ts.net" || true
      if ! networksetup -setdnsservers "$ACTIVE_SERVICE" "$target_ip"; then
        echo "FAILED (networksetup error  check permissions)"
        sleep 1
        continue
      fi
      networksetup -setsearchdomains "$EN0_SERVICE" "ts.net" || true
      networksetup -setdnsservers "$EN0_SERVICE" "$target_ip" || true

      local entries
      entries=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log \
                | awk '/^(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)
    \.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)$/ {print}')
      if echo "$entries" | grep -qFx "$target_ip"; then
        echo "VERIFIED"
        return 0
      fi

      local current
      current=$(echo "$entries" | tr '\n' ' ' | sed 's/ $//')
      echo "NOT YET (current DNS: [${current:-none}])"

      if [ "$attempt" = "3" ]; then sleep 2; else sleep 1; fi
    done

    echo "  [force_dns] FAILED after 3 attempts"
    return 1
  else
    local target_ip="$1"

    for attempt in {1..3}; do
      echo -n "  [force_dns] Attempt $attempt/3: setting DNS to $target_ip on $ACTIVE_SERVICE... "
      sudo -n resolvectl domain "$ACTIVE_SERVICE" "~ts.net" 2>/dev/null || true
      if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" "$target_ip"; then
        echo "FAILED (resolvectl error)"
      else
        # Verify
        entries=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+')
        if [ "$entries" == "$target_ip" ]; then
          echo "VERIFIED"
          return 0
        else
          echo "MISMATCH (Resolver refused lock)"
        fi
      fi
      sleep 2
    done
    return 1
  fi
}
```

12 scripts/watcher.sh

```
#!/bin/bash
# scripts/watcher.sh  nullexit post-wake / post-roam recovery watcher
#
# Long-running daemon. Launched by launchd as a LaunchAgent at
# ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
# (label com.nullexit.wake-recovery).
```

```
#
# Two independent listeners run as backgrounded pipelines:
#
#   (1) WAKE listener  `log stream` live-follows the macOS unified log for
#       power-management events (lid closewake, manual wake, DarkWake from
#       clamshell). Each event triggers run_recover (with a global debounce so
#       a single wake that emits 2-3 log lines only causes ONE post-wake run).
#
#   (2) NETWORK listener  `scutil n.watch` on `State:/Network/Global/IPv4`
#       fires on every Wi-Fi roam, ethernet swap, hotspot join, captive-portal
#       re-bind, VPN up/down, etc. Same debounce.
#
# Both listeners call run_recover, which:
#   - checks /tmp/nullexit-gateway-active.marker (only act when toggle.sh left
#     the gateway up)
#   - debounces (10s default) so multiple events don't pile up
#   - shells out to `bash <repo>/recover.sh --post-wake` directly (no GUI auth prompt needed due to NOPASSWD
    sudoers)
#
# See devref.md 10.29 for the full Apple power management primer and the
# rationale behind using `log stream` + `scutil n.watch` from a shell daemon.

# NOTE: errexit (`set -e`) is intentionally NOT enabled. This is a long-running
# daemon, not a discrete-step script, and errexit actively breaks the reconnect
# loops below: any failed pipe (log stream on rotation, scutil on state churn)
# would kill the outer subshell before `echo "reconnecting..."; sleep 5;` runs,
# AND the parent's final `wait` would propagate a non-zero child exit and kill
# the entire watcher. Every error path is already wrapped in `|| true` /
# `2>/dev/null` fallbacks.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
source "$SCRIPT_DIR/common.sh"
setup_standard_path

LOG="$SCRIPT_DIR/../output.log"
# Resolve our own directory; recover.sh lives one level up at the repo root.
# This makes the daemon install-agnostic  no need to template the path
# in launchd, no need to keep a deploy-specific hardcoded location in sync.
RECOVER="$SCRIPT_DIR/../recover.sh"
MARKER="$MARKER_FILE"
DEBOUNCE_FILE="/tmp/nullexit-watcher.last-recovery"
DEBOUNCE_SECONDS="${NULLEXIT_DEBOUNCE_SECONDS:-10}"
# Post-recovery SETTLE cooldown. A post-wake run re-asserts the exit node
# (recover.sh: `tailscale set --exit-node`), which re-writes the host default
# route  and that route IS State:/Network/Global/IPv4, the exact key Listener 2
# watches. The self-triggered change lands ~15-45s AFTER the run finishes, past
# the 10s debounce, so without a longer window the watcher re-fires post-wake in
# a ~40s self-feedback loop that only stops once routing happens to reach a fixed
# point (devref 15.12.21). This cooldown is armed ONLY after a recovery runs, so
# a genuine first wake/roam is still handled immediately; only the change we
# caused ourselves gets absorbed. A real new roam during the window is delayed at
# most this many seconds, which is acceptable.
COOLDOWN_FILE="/tmp/nullexit-watcher.settle-until"
POSTWAKE_SETTLE_SECONDS="${NULLEXIT_POSTWAKE_SETTLE_SECONDS:-45}"

exec >> "$LOG" 2>&1

#  Single-instance lock
# Without this, repeated lid-closewake cycles under macOS Sonoma+ accumulate
# orphan watcher.sh processes: launchd suspends on sleep and SIGCONT-resumes
# on wake, and KeepAlive.Crashed=true caused relaunch on any non-zero exit;
# the listener grandchildren (log stream | while ...  and  (echo n.add ...;
# sleep 86400) | scutil | while ...) stay alive because pipe producers don't
# reliably respond to plain SIGTERM if the consumer is in a half-closed state.
#
# We solve it with a PID-file lock (POSIX-portable; macOS does NOT ship `flock`).
# The trap below removes the lock file on every exit so a stale lock never
# outlives a crashed watcher. On launch: read the existing PID, check it is
# alive AND is actually a `scripts/watcher.sh` process; if so, exit 0; else
# overwrite with our PID and proceed. The check-then-write window has a TOCTOU
# race but is microseconds wide and only matters for hand-launched duplicate
# test invocations; production is single-launch via launchctl load -w.
LOCK_FILE="/tmp/nullexit-watcher.lock"
LOCK_PID=""
if [ -e "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    # Verify the holder is actually a scripts/watcher.sh process (not some
    # unrelated process that got the recycled PID). The exact-launchd pattern
    # `bash <abs-path>/scripts/watcher.sh` is hard to accidentally match with
    # a `grep scripts/watcher.sh .` invocation because we anchor on `^bash`
    # followed by whitespace and end-of-line on the script path.
```

```
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -qE "(^|[[:space:]])bash[[:space:]]+.*scripts/watcher\.sh$
     "; then
      echo "[$(date -u +%FT%TZ)] watcher.sh: another instance holds $LOCK_FILE (pid=$LOCK_PID), exiting"
      exit 0
    fi
  fi
fi
echo $$ > "$LOCK_FILE"

echo "[$(date -u +%FT%TZ)] watcher.sh start (pid=$$ ppid=$PPID user=$(id -un) lock=acquired)"

# Treat launchd-resume as a wake signal. Apple's launchd SUSPENDS LaunchAgents
# during system sleep and restores them on wake; during that suspended gap,
# macOS's `log stream` cannot catch the wake event that prompted the resume.
# The next thing that happens after wake IS our own startup, so we always
# fire one post-wake run unconditionally here. The marker check + debounce
# in run_recover() make this idempotent: if the gateway isn't up or we just
# debounced, it no-ops. Without this, the FIRST wake-after-install goes
# unseen and the watcher appears broken until the SECOND wake.
#  Helper: run recover.sh --post-wake if the gateway is active
run_recover() {
  local why="$1"
  if [ ! -f "$MARKER" ]; then
    echo "[$(date -u +%FT%TZ)] $why  no $MARKER, skip"
    return 0
  fi
  local now last settle_until
  now=$(date +%s)
  # Post-recovery settle window: ignore the Global/IPv4 change our own last
  # post-wake run caused (exit-node re-assert) so we don't self-feedback loop.
  settle_until=$(cat "$COOLDOWN_FILE" 2>/dev/null || echo 0)
  if [ "$now" -lt "$settle_until" ]; then
    echo "[$(date -u +%FT%TZ)] $why  settling ($((settle_until - now))s left in post-recovery cooldown), skip"
    return 0
  fi
  last=$(cat "$DEBOUNCE_FILE" 2>/dev/null || echo 0)
  if [ "$((now - last))" -lt "$DEBOUNCE_SECONDS" ]; then
    echo "[$(date -u +%FT%TZ)] $why  debounced ($((now - last))s since last, threshold ${DEBOUNCE_SECONDS}s)"
    return 0
  fi
  echo "$now" > "$DEBOUNCE_FILE"
  echo "[$(date -u +%FT%TZ)] $why  bash $RECOVER --post-wake"
  bash "$RECOVER" --post-wake
  local exit_code=$?
  now=$(date +%s)
  echo "$now" > "$DEBOUNCE_FILE"
  # Arm the settle cooldown so the route change this run just caused can't re-fire us.
  echo "$((now + POSTWAKE_SETTLE_SECONDS))" > "$COOLDOWN_FILE"
  echo "[$(date -u +%FT%TZ)] $why  exit=$exit_code (settle until +${POSTWAKE_SETTLE_SECONDS}s, now=$now)"
  return $exit_code
}

if [ -f "$MARKER" ]; then
  run_recover "initial post-wake (launchd resume = wake signal)"
fi

#  LAN P2P Auto-Detection
# Detects whether the current network allows safe Tailscale LAN P2P by:
#   1. Checking WPA2-Enterprise (802.1x) via system_profiler SPAirPortDataType  always forces P2P=false
#      because enterprise networks almost universally have AP Client Isolation.
#   2. AP Isolation probe  sends a broadcast ping and checks if ANY LAN host responds.
#      If no LAN hosts respond on a non-empty network, AP Isolation is active.
# Writes 'true' or 'false' to .lan_p2p_detected in the repo root.
# routing-fix.sh reads this file dynamically every 30s loop iteration.
P2P_OVERRIDE_FILE="$SCRIPT_DIR/../.lan_p2p_detected"


detect_lan_p2p_mode() {
  local reason="unknown" allow="false"

  # Step 1: Detect Wi-Fi security type via system_profiler (airport binary removed in macOS 14.4+)
  # SPAirPortDataType lists all known networks; the first "Security:" line is the current association.
  local security
  security=$(system_profiler SPAirPortDataType 2>/dev/null \
    | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF; exit}')

  if [ -z "$security" ]; then
    # Not on Wi-Fi or system_profiler unavailable  fail safe
    reason="no-wifi" allow="false"
  elif echo "$security" | grep -qi "Enterprise"; then
    # WPA2/WPA3 Enterprise = 802.1x  always AP isolated
```

```
    reason="802.1x-enterprise" allow="false"
  elif echo "$security" | grep -qi "Personal\|None\|Open"; then
    # WPA2/WPA3 Personal or open  probe for AP isolation
    local gw
    gw=$(route -n get 1.1.1.1 2>/dev/null | awk '/gateway:/{print $2}')
    if [ -z "$gw" ]; then
      reason="no-default-route" allow="false"
    else
      local responses
      responses=$(ping -c 3 -t 1 -q "$gw" 2>/dev/null | awk '/packets transmitted/{print $4}')
      if [ "${responses:-0}" -gt 0 ]; then
        reason="wpa-personal-lan-reachable" allow="true"
      else
        reason="wpa-personal-ap-isolated" allow="false"
      fi
    fi
  else
    # Unknown security type  fail safe
    reason="unknown-security-type" allow="false"
  fi

  echo "[$(date -u +%FT%TZ)] P2P detect: security='${security:-none}'  allow=${allow} (reason: ${reason})"
  echo "$allow" > "$P2P_OVERRIDE_FILE"
  write_host_ips
}

# Run once on startup
detect_lan_p2p_mode

#  Listener 1: WAKE events via unified log
# Predicate picks up:
#   * "didWake"              regular kernel wake events (lid open, RTC alarm)
#   * "Wake from Sleep"      powerd-driven normal wake
#   * "DarkWake"             clamshell-display wake that didn't fully wake the
#                             Mac (relevant when lid opens during display sleep)
#   * "Waking from"          generic catch-all for variations in newer macOS
# `[c]` makes the match case-insensitive (the unified log uses mixed-case
# phrases). --info reduces noise by dropping debug/trace log levels.
# Subsystem-based predicate catches every powermanagement emission (willSleep,
# didWake, didDim, didUndim, DarkWake, etc.) regardless of message phrasing
# across kernel versions. Phrasing-based predicates ('didWake', 'Wake from Sleep')
# are fragile across macOS releases. We accept the extra log noise because the
# run_recover() debounce + marker check filter out anything that isn't a real
# gateway-relevant event.
(
  # Reconnect loop: macOS rotates the unified log buffer periodically; when
  # that happens, this `log stream` subscriber's pipe EOFs, the inner
  # `while read` consumer terminates, and without this wrapper the listener
  # would be silently dead for the rest of the watcher's lifetime.
  while :; do
    log stream --predicate \
      'subsystem == "com.apple.powermanagement"' \
       --style compact --info 2>/dev/null \
      | while IFS= read -r line; do
          [ -z "$line" ] && continue
          run_recover "WAKE: $(echo "$line" | head -c 100)"
        done
    echo "[$(date -u +%FT%TZ)] log stream subshell exited; reconnecting in 5s"
    sleep 5
  done
) &

#  Listener 2: NETWORK state changes via scutil
# `scutil n.watch` on a state key prints SCEventUpdate lines whenever the key
# changes. We watch Global/IPv4 because that's the umbrella that catches link,
# address, route, and DNS changes for ALL interfaces in one namespace.
# The pipe is kept alive by an `sleep 86400` loop so scutil never sees EOF
# from us (which would silently terminate the watch).
(
  # Reconnect loop: scutil can exit noisily on certain state changes (rare,
  # but observed). Without the wrapper, the network listener's pipe EOFs and
  # every subsequent roam goes unseen until launchd restarts us.
  while :; do
    (
      # Watch both IPv4 and IPv6 umbrella state  IPv6-only or dual-stack
      # networks never fire on the IPv4 key alone, and the user's first roam
      # in a hotel / airport / on cellular hotspot is often IPv6-first under
      # NAT64.
      echo "n.add State:/Network/Global/IPv4"
      echo "n.add State:/Network/Global/IPv6"
      echo "n.watch"
```

```
      while true; do sleep 86400; done
    ) | script -q /dev/null scutil 2>/dev/null \
      | while IFS= read -r line; do
          case "$line" in
            *changedKey*|*State:/Network/Global/IPv*|*n.state*|*SCEventUpdate*)
              detect_lan_p2p_mode
              run_recover "NET: $(echo "$line" | head -c 100)"
              ;;
            *) ;;
          esac
        done
    echo "[$(date -u +%FT%TZ)] scutil pipe exited; reconnecting in 5s"
    sleep 5
  done
) &
#  Listener 3: TUNNEL_FAILED_CLOSED.marker Watcher
# Polls for the fail-closed marker dropped by the routing-fix container.
(
  last_notified=0
  while :; do
    if [ -f "$SCRIPT_DIR/../TUNNEL_FAILED_CLOSED.marker" ]; then
      now=$(date +%s)
      # Notify once every 5 minutes (300s) if it stays failed
      if [ "$((now - last_notified))" -ge 300 ]; then
        osascript -e 'display notification "VPN Tunnel Health Check Failed. Internet traffic is blackholed to prevent IP leak." with title "NullExit Alert"' || true
        last_notified=$now
      fi
    else
      last_notified=0
    fi
    sleep 3
  done
) &

#  Listener 4: Wi-Fi uplink keepalive
# Re-associates a remembered network (scripts/wifi-rejoin.sh) whenever the
# Wi-Fi interface has no routable IP. This lives HERE, in the always-on
# wake-recovery daemon, on purpose: the WARP watcher (toggle.sh) that used to
# be the ONLY driver of wifi-rejoin runs solely while the gateway is UP, so a
# Wi-Fi drop while the gateway was down/off left the Mac "just sitting there
# without internet"  nothing rejoined. This poller is independent of gateway
# state, so it recovers Wi-Fi either way. Leak-safe: it only acts when the
# interface has NO address (empty or 169.254.* self-assigned)  i.e. no egress
# path exists, so there is nothing to leak. wifi-rejoin.sh enforces its own
# grace period, backoff, and give-up cap, and only ever joins a network you
# explicitly remembered. When the gateway IS up, the WARP watcher may also call
# wifi-rejoin; the two share the same /tmp markers so they coordinate rather
# than double-associate.
(
  WIFI_HELPER="$SCRIPT_DIR/wifi-rejoin.sh"
  wifi_iface=$(networksetup -listallhardwareports 2>/dev/null \
    | awk '/Hardware Port: Wi-Fi/{getline; print $2; exit}')
  [ -z "$wifi_iface" ] && wifi_iface=en0
  while :; do
    phys_ip=$(ipconfig getifaddr "$wifi_iface" 2>/dev/null || true)
    case "${phys_ip:-}" in
      ''|169.254.*) bash "$WIFI_HELPER" rejoin >/dev/null 2>&1 || true ;;
      *)            bash "$WIFI_HELPER" reset  >/dev/null 2>&1 || true ;;
    esac
    sleep 15
  done
) &

#  Lifetime
# `wait` blocks until both background pipelines exit (which they normally
# never do). launchctl `bootout`/SIGTERM triggers the trap, which kills the
# whole process group so the sleep-forever in the scutil pipe also dies.
# We catch TERM/INT/HUP for explicit signals and EXIT for any normal-exit path
# (so listeners always get cleaned up even if `exit 0` runs somewhere).
# `rm -f $LOCK_FILE` first so a stale lock can never block the next launch.
# `pkill -TERM -P $$` then targets direct listener-children, then `kill 0`
# takes care of any backgrounded group-mates not reachable via the parent.
trap 'echo "[$(date -u +%FT%TZ)] watcher.sh terminating (signal/exit=$?)"; rm -f "$LOCK_FILE" 2>/dev/null || true; pkill -TERM -P $$ 2>/dev/null || true; kill 0 2>/dev/null || true; exit 0' TERM INT HUP EXIT
wait
```

13 scripts/crypto.sh

```
#!/usr/bin/env bash
# crypto.sh - Cryptographic Integrity Check

set -euo pipefail
cd "$(dirname "$0")/.."

if [ "$#" -ne 1 ] || { [ "$1" != "--sign" ] && [ "$1" != "--verify" ]; }; then
  echo "Usage: $0 --sign | --verify"
  exit 1
fi

if [ ! -f .env ]; then
  echo "Error: .env not found."
  exit 1
fi

SEED=$(grep "^NULLEXIT_SEED=" .env | cut -d '=' -f2)
if [ -z "$SEED" ]; then
  echo "Error: NULLEXIT_SEED not found in .env."
  exit 1
fi

if [ "$1" == "--sign" ]; then
  echo "Generating cryptographic signatures using seed..."
  rm -f .signatures
  for file in toggle.sh recover.sh scripts/common.sh scripts/routing-fix.sh scripts/diagnose-host-leak.sh
     scripts/sweep.sh scripts/noise.sh scripts/watcher.sh scripts/pf.conf post-rules.txt docker-compose.yml
     black_list.txt white_list.txt; do
    if [ -f "$file" ]; then
      hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
      echo "$file:$hash" >> .signatures
      echo "  [Signed] $file"
    else
      echo "  [Skipped] $file not found"
    fi
  done
  echo "Signatures saved to .signatures"
  exit 0
fi

if [ "$1" == "--verify" ]; then
  if [ ! -f .signatures ]; then
    echo "Error: .signatures file missing. Run $0 --sign first."
    exit 1
  fi
  FAIL=0
  while IFS=: read -r file expected_hash; do
    if [ -z "$file" ]; then continue; fi
    if [ ! -f "$file" ]; then
      echo "[FAIL] $file is missing!"
      FAIL=1
      continue
    fi
    actual_hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
    if [ "$actual_hash" != "$expected_hash" ]; then
      echo "[FAIL] $file has been modified! Hash mismatch."
      FAIL=1
    fi
  done < .signatures

  if [ "$FAIL" -eq 1 ]; then
    echo ""
    echo "CRITICAL: Script integrity verification failed!"
    echo "One or more core files have been modified without authorization."
    echo "To authorize these changes, run: ./scripts/crypto.sh --sign"
    echo ""
    exit 1
  fi
  exit 0
fi
```

14 .aiignore

```
# Environment variables and secrets
.env
tor-bridges.txt
```

15 .claude/settings.local.json

```
{
  "permissions": {
    "allow": [
      "Bash(grep *)",
      "Bash(python3 /tmp/gen_nullexit_graph.py)",
      "Bash(python3 -)",
      "Bash(\"/Applications/Google Chrome.app/Contents/MacOS/Google Chrome\" --headless=new --disable-gpu --window-size=1280,800 --virtual-time-budget=8000 --dump-dom \"file:///tmp/kg_debug.html\")",
      "Bash(echo \"dangling-check-exit=$?\")",
      "Bash(bash scripts/wifi-rejoin.sh reset)",
      "Bash(launchctl kickstart *)",
      "Bash(perl -0pi -e 's/see 15\\\\.11\\\\.7\\\\b/see 15.11.10/g; s/\\\\\\(15\\\\.11\\\\.7\\\\\\)/\\(15.11.10\\)/g; s/15\\\\.11\\\\.7 \\\\\\(SSID redaction\\\\\\)/15.11.10 \\(SSID redaction\\)/g' devref.md)",
      "Bash(perl -0pi -e 's/see `devref\\\\.md 15\\\\.11\\\\.7`/see `devref.md 15.11.10`/g' README.md)"
    ]
  }
}
```

16 .env.example

```
# nullexit configuration template  copy to .env and fill in.
#   cp .env.example .env
# .env is gitignored (it holds secrets); this template documents every option
# with its shipping default. For coherent bundles, see scripts/profiles.sh.

#  Secrets (fill these in)
# WARP WireGuard keys  from your Cloudflare WARP registration.
WIREGUARD_PRIVATE_KEY=<your-warp-wireguard-private-key>
WIREGUARD_PUBLIC_KEY=<your-warp-wireguard-public-key>
WIREGUARD_ADDRESSES=172.16.0.2/32
# Tailscale auth key  Admin Console  Settings  Keys  Generate auth key.
TS_AUTHKEY=tskey-auth-xxxxxxxxxxxxxxxxxxxxxxx
# HMAC seed for script-integrity signatures  generate with: openssl rand -hex 32
NULLEXIT_SEED=<run: openssl rand -hex 32>
# AdGuard Home UI + Tor ControlPort credentials.
ADGUARD_USER=admin
ADGUARD_PASSWORD=<choose-a-password>

#  Gateway core
# Rule-compiler memory profile: light (~52k) / medium (~167k) / heavy (~253k).
GATEWAY_RULE_PROFILE=medium
STOP_COLIMA_ON_EXIT=true
# PF fail-closed kill-switch. Leave true  this is the core leak guarantee.
KILL_SWITCH=true
# true only on trusted home networks (no AP isolation); false on campus/enterprise
# WPA2-Enterprise (forces DERP, avoids gvproxy SNAT poisoning).
TAILSCALE_ALLOW_LAN_P2P=false
GATEWAY_MSS=1120
GATEWAY_BYPASS_PING=false
GATEWAY_USE_EXIT_NODE=true
# Consecutive warp=off polls before auto-shutdown (6 = 30s; 3 = faster/stricter).
WARP_FAIL_THRESHOLD=6
BLOCKED_COUNTRIES="kp il"
HOST_MTU=1200
# Debug xtrace of toggle.sh to output.log (keep false for normal use).
DEBUG_TRACE=false

#  Tor hardening
# true = Tor must use obfs4 bridges only (never public guards).
TOR_USE_BRIDGES=false
TOR_BRIDGE_LINES_FILE=./tor-bridges.txt
TOR_EXCLUDE_EXIT_NODES={us},{gb},{fr},{de},{it},{ca},{jp},{kp},{il}

#  Cover traffic / dummy padding (scripts/noise.sh)  opt-in, default off
```

```
NOISE_ENABLED=false
NOISE_PAD_TARGET=
NOISE_PAD_PORT=9999
NOISE_PAD_RATE_KBPS=64
NOISE_PAD_PACKET_BYTES=1024
NOISE_PAD_JITTER_MS=40
# constant = always emit the target rate; topup = fill real-traffic gaps to keep
# the observed rate flat (stronger anti-correlation, no waste when already busy).
NOISE_PAD_MODE=constant
NOISE_PAD_IFACE=en0
# Optional: target as a % of LINK CAPACITY (overrides NOISE_PAD_RATE_KBPS).
NOISE_PAD_PCT=
NOISE_LINK_MBPS=100

#  FaceTime split-route (scripts/common.sh)  opt-in, default off
#  true routes the Apple /16s below DIRECT on your REAL IP (not WARP) so
# FaceTime can hole-punch. A deliberate, narrow hole scoped to Apple infra.
FACETIME_SPLIT_ROUTE=false
FACETIME_DIRECT_SUBNETS="17.249.0.0/16"

#  LAN Pi-hole mode (scripts/pihole-mode.sh)  opt-in, default off
# Serve AdGuard DNS filtering to non-Tailscale LAN devices (trusted networks only).
PIHOLE_LAN_MODE=false
PIHOLE_LAN_LISTEN=

#  Sweep throughput check (scripts/sweep.sh, check 7/7)  all optional
# Uncomment to override the defaults baked into sweep.sh (a fast CDN, 30s cap).
# SWEEP_THROUGHPUT_URL=https://ash-speed.hetzner.com/100MB.bin
# SWEEP_THROUGHPUT_TIMEOUT=30
# SWEEP_SKIP_THROUGHPUT=false
```

17 .gitignore

```
# Environment variables and secrets
.env

# wgcf generated files containing keys and account info
wgcf-account.toml
wgcf-profile.conf

# Tailscale persistent state containing node identity
tailscale-state/

# Binaries
wgcf

# macOS App bundles
*.app/
*.app

# AdGuard Home persistent state
adguard/
.gateway_ip
logs.txt

#
stability_incident_report.md
output.log
.DS_Store
*.desktop
blocked.log
wtf.md
host-leak-diagnostic-*.txt
nullexit_review.md
nullexit_unified.tex
nullexit_unified.txt

nullexit_unified.synctex.gz
tor-bridges.txt

# Runtime state files
.host_ips
.lan_p2p_detected
.build_hash
.rules_hash
.last_wifi_ssid
```

```
.last_wifi_ssid.tmp
.signatures
TUNNEL_FAILED_CLOSED.marker

# Python cache
__pycache__/
*.py[cod]

# LaTeX build artifacts
*.aux
*.log
*.out
*.toc
*.xdv
*.fdb_latexmk
*.fls
*.synctex*
refactor.md
sweep.md

# Obsidian Knowledge graph
nullexit-knowledge-graph/

# DNS anomaly detector  per-network model (regenerated by 'learn')
.dns_baseline.json

# .env backups written by profiles.sh (contain secrets  never commit)
.env.bak*

# device-scramble saved real identity (never commit)
.device-identity.orig
```

18 Recover-Gateway.applescript

```
tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
try
  do shell script "true" with administrator privileges
on error
  return
end try
tell application "Terminal"
  activate
  do script "cd " & quoted form of scriptDir & "; bash recover.sh"
end tell
```

19 Sweep-Gateway.applescript

```
tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
try
  do shell script "cd " & quoted form of scriptDir & " && ./scripts/crypto.sh --verify"
on error errMsg
  display dialog "nullexit integrity check FAILED  sweep will not run." & return & return & "A signed file (
     scripts/sweep.sh, scripts/common.sh, ) no longer matches .signatures. If you edited it on purpose, re-sign
     with:" & return & "./scripts/crypto.sh --sign" & return & return & errMsg buttons {"Abort"} default button
     "Abort" with icon stop
  return
end try
tell application "Terminal"
  activate
  do script "cd " & quoted form of scriptDir & " && bash scripts/sweep.sh"
end tell
```

20 Toggle-Gateway.applescript

```
tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
try
  do shell script "true" with administrator privileges
on error
```

```
  return
end try
tell application "Terminal"
  activate
  do script "cd " & quoted form of scriptDir & "; ./toggle.sh"
end tell
```

21 benchmarks/benchmark.sh

```
#!/usr/bin/env bash

# benchmark.sh
# Automates the full benchmarking suite (both Python download test and Lighthouse HTML test)
# Runs the suite with Gateway ON, toggles OFF, runs again, then toggles back ON.

set -e

# Change to project root so we can access toggle.sh reliably
cd "$(dirname "$0")/.."

# Ensure jq is installed
if ! command -v jq &> /dev/null; then
    echo "Error: 'jq' is not installed. Please run 'brew install jq'"
    exit 1
fi

run_lighthouse() {
    local url=$1
    local file_prefix=$2
    echo "  [Lighthouse] Testing $url ..."

    # Run lighthouse headlessly, silencing standard output
    npx -y lighthouse "$url" --chrome-flags="--headless" --only-categories=performance --output=json --output-
     path="${file_prefix}.json" >/dev/null 2>&1

    # Parse the results
    local score=$(jq -r '.categories.performance.score' "${file_prefix}.json")
    local tti=$(jq -r '.audits["interactive"].displayValue' "${file_prefix}.json")
    local si=$(jq -r '.audits["speed-index"].displayValue' "${file_prefix}.json")

    # Convert score to percentage
    local score_pct=$(awk "BEGIN {print $score*100}")

    echo "    -> Performance Score: ${score_pct}%"
    echo "    -> Time to Interactive: $tti"
    echo "    -> Speed Index: $si"

    # Cleanup
    rm -f "${file_prefix}.json"
}

echo "============================================="
echo "   PHASE 1: Testing with Gateway ON"
echo "============================================="
# 1. Run raw Python download test
python3 benchmarks/test_load.py

# 2. Run Lighthouse HTML render test
echo ""
echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
run_lighthouse "https://www.cnet.com/" "on_cnet"
run_lighthouse "https://www.independent.co.uk/" "on_ind"


echo ""
echo "============================================="
echo "   PHASE 2: Toggling Gateway OFF"
echo "============================================="
./toggle.sh
echo "Waiting 10 seconds for macOS Wi-Fi to stabilize after DNS reset..."
sleep 10


echo ""
echo "============================================="
echo "   PHASE 3: Testing with Gateway OFF"
```

```
echo "============================================="
# 1. Run raw Python download test
python3 benchmarks/test_load.py

# 2. Run Lighthouse HTML render test
echo ""
echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
run_lighthouse "https://www.cnet.com/" "off_cnet"
run_lighthouse "https://www.independent.co.uk/" "off_ind"


echo ""
echo "============================================="
echo "   PHASE 4: Restoring Gateway ON"
echo "============================================="
./toggle.sh

echo ""
echo "============================================="
echo " BENCHMARKS COMPLETE!"
echo "You can scroll up to compare Phase 1 (AdGuard ON) vs Phase 3 (AdGuard OFF)."
echo "============================================="
```

22 benchmarks/test_load.py

```
import urllib.request
import urllib.error
import time
from html.parser import HTMLParser
import concurrent.futures

class ResourceParser(HTMLParser):
    def __init__(self):
        super().__init__()
        self.urls = set()

    def handle_starttag(self, tag, attrs):
        attrs = dict(attrs)
        url = None
        if tag in ['img', 'script']:
            url = attrs.get('src')
        elif tag == 'link' and attrs.get('rel') == 'stylesheet':
            url = attrs.get('href')

        if url and url.startswith('http'):
            self.urls.add(url)

def fetch_url(url):
    try:
        req = urllib.request.Request(url, headers={'User-Agent': 'Mozilla/5.0'})
        urllib.request.urlopen(req, timeout=3)
        return True
    except Exception:
        return False

def measure_url(target_url):
    print(f"\nTesting {target_url} Load...")
    start_total = time.time()

    # Fetch main HTML
    try:
        req = urllib.request.Request(target_url, headers={'User-Agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X
      10_15_7) AppleWebKit/537.36'})
        response = urllib.request.urlopen(req, timeout=5)
        html = response.read().decode('utf-8', errors='ignore')
    except Exception as e:
        print(f"Failed to fetch {target_url}: {e}")
        return

    parser = ResourceParser()
    parser.feed(html)
    urls = list(parser.urls)
    print(f"Found {len(urls)} subresources to fetch (scripts, images, css)...")

    # Concurrently fetch subresources
    success_count = 0
```

```
    fail_count = 0
    max_threads = min(64, max(1, len(urls)))
    with concurrent.futures.ThreadPoolExecutor(max_workers=max_threads) as executor:
        results = list(executor.map(fetch_url, urls))

    success_count = sum(1 for r in results if r)
    fail_count = len(results) - success_count

    total_time = time.time() - start_total
    print(f"Time taken: {total_time:.2f} seconds")
    print(f"Successfully fetched: {success_count}, Failed/Blocked: {fail_count}")

urls_to_test = [
    "https://www.dailymail.co.uk/home/index.html",
    "https://www.cnet.com/",
    "https://www.independent.co.uk/"
]

for u in urls_to_test:
    measure_url(u)
```

23 black_list.txt

```
ad.doubleclick.net
adclick.g.doubleclick.net
adeventtracker.spotify.com
adfstat.yandex.ru
ads-api.tiktok.com
ads-api.twitter.com
ads-fa.spotify.com
ads-sg.tiktok.com
ads.tiktok.com
adserver.unityads.unity3d.com
adsfs.oppomobile.com
an.facebook.com$important
connect.facebook.net$important
pixel.facebook.com$important
analytics.query.yahoo.com
analytics.s3.amazonaws.com
analytics.spotify.com
analyticsengine.s3.amazonaws.com
appmetrica.yandex.ru
audio2.spotify.com
bdapi-ads.realmemobile.com
bdapi-in-ads.realmemobile.com
bounceexchange.com
bs.serving-sys.com
business-api.tiktok.com
ck.ads.oppomobile.com
crashdump.spotify.com
data.ads.oppomobile.com
data.mistat.india.xiaomi.com
data.mistat.rus.xiaomi.com
data.mistat.xiaomi.com
desktop.spotify.com
doubleclick.net
ds.metric.spotify.com
gemini.yahoo.com
googleads.g.doubleclick.net
googleadservices.com
googlesyndication.com
grs.hicloud.com
gtm.spotify.com
iot-eu-logser.realme.com
iot-logser.realme.com
log.byteoversea.com
log.fc.yahoo.com
log.spotify.com
m.doubleclick.net
mediavisor.doubleclick.net
metrics.spotify.com
metrika.yandex.ru
omaze.com
pagead2.googlesyndication.com
securepubads.g.doubleclick.net
static.doubleclick.net
```

```
stats.g.doubleclick.net
tracking.rus.miui.com
udcm.yahoo.com
v.jwpcdn.com
weblb-wg.gslb.spotify.com
webview.unityads.unity3d.com
diagassets.apple.com
iphonesubmissions.apple.com
radarsubmissions.apple.com
metrics.apple.com
xp.apple.com
heads-fa.spotify.com
heads4.spotify.com
pixel.spotify.com
segs.spotify.com
audio-fa.spotify.com
events.data.microsoft.com
datarouter-prod.ak.epicgames.com
nel.cloudflare.com
eu-mobile.events.data.microsoft.com
stocks-analytics-events.apple.com
settings-win.data.microsoft.com
choice.microsoft.com
watson.live.com$important
gdmf.apple.com
deviceenrollment.apple.com
deviceservices-external.apple.com
gpm.samsungqbe.com
daisy.ubuntu.com
popcon.ubuntu.com
popcon.debian.org
mdm-enrollment.googleapis.com$important
device-provisioning.googleapis.com$important
```

24 cloud-init.yaml

```yaml
#cloud-config
# Generic Cloud-Init Bootstrap Script for Remote Exit Nodes (Ubuntu/Debian)
# Paste this into the "User Data" or "Initialization Script" section when creating a VPS.

package_update: true
package_upgrade: true

packages:
  - apt-transport-https
  - ca-certificates
  - curl
  - gnupg
  - lsb-release
  - git
  - python3
  - python3-pip

runcmd:
  # 1. Install Docker natively
  - curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg
  - echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
  - sudo apt-get update
  - sudo apt-get install -y docker-ce docker-ce-cli containerd.io docker-compose-plugin

  # 2. Add default user (ubuntu) to docker group
  - usermod -aG docker ubuntu || true
  - usermod -aG docker debian || true

  # 3. Create the installation directory
  - mkdir -p /opt/nullexit
  - chown -R ubuntu:ubuntu /opt/nullexit || chown -R debian:debian /opt/nullexit

  # 4. Enable IPv4 and IPv6 forwarding for VPN routing
  - echo "net.ipv4.ip_forward=1" >> /etc/sysctl.d/99-tailscale.conf
  - echo "net.ipv6.conf.all.forwarding=1" >> /etc/sysctl.d/99-tailscale.conf
  - sysctl -p /etc/sysctl.d/99-tailscale.conf

```

```
final_message: "The remote exit node has been fully bootstrapped. SSH in, clone the repo to /opt/nullexit, add your .env, and run ./setup-linux.sh!"
```

25 docker/tor/Dockerfile

```
FROM debian:bookworm-slim

RUN apt-get update && \
    apt-get install -y tor obfs4proxy gosu bash && \
    rm -rf /var/lib/apt/lists/*

COPY entrypoint.sh /entrypoint.sh
RUN chmod +x /entrypoint.sh

ENTRYPOINT ["/entrypoint.sh"]
```

26 docker/tor/entrypoint.sh

```
#!/bin/bash
set -e

TORRC="/etc/tor/torrc"
mkdir -p /etc/tor /var/lib/tor

# Base configuration: SOCKS 9050, Control 9051, TransPort 9040, DNSPort 5353  all bound to
# 0.0.0.0 (required for Docker/Colima host port-mapping). The ControlPort is secured via a
# dynamic HashedControlPassword (added below when TOR_PASSWORD is set); see devref.md 15.10.2.
cat <<EOF > $TORRC
SocksPort 0.0.0.0:9050
ControlPort 0.0.0.0:9051
CookieAuthentication 0
Log notice stdout
DataDirectory /var/lib/tor

# Transparent proxying & DNS resolution for .onion mapping
TransPort 0.0.0.0:9040
DNSPort 0.0.0.0:5353
AutomapHostsOnResolve 1
VirtualAddrNetworkIPv4 198.18.0.0/15
EOF

if [ -n "$TOR_PASSWORD" ]; then
    HASH=$(tor --hash-password "$TOR_PASSWORD" | tail -n 1)
    echo "HashedControlPassword $HASH" >> $TORRC
fi

if [ -n "$TOR_EXCLUDE_EXIT_NODES" ]; then
    echo "ExcludeExitNodes $TOR_EXCLUDE_EXIT_NODES" >> $TORRC
    echo "StrictNodes 1" >> $TORRC
fi

if [ "${TOR_USE_BRIDGES:-false}" = "true" ]; then
    BRIDGE_FILE="${TOR_BRIDGE_LINES_FILE:-/tor-bridges.txt}"

    if ! grep -vE "^[[:space:]]*#" "$BRIDGE_FILE" | grep -q '[^[:space:]]'; then
        MSG="[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL] [Tor Proxy] TOR_USE_BRIDGES is enabled, but no bridge lines found at $BRIDGE_FILE. Proxy startup ABORTED. Get bridges from https://bridges.torproject.org."
        echo "$MSG"
        if [ -w "/output.log" ]; then
            echo "$MSG" >> /output.log
        fi
        # Sleep indefinitely to prevent docker-compose 'unless-stopped' from triggering a crash loop
        exec sleep infinity
    fi

    echo "UseBridges 1" >> $TORRC
    echo "ClientTransportPlugin obfs4 exec /usr/bin/obfs4proxy" >> $TORRC

    while IFS= read -r line; do
        # Ignore comments and empty lines
        [[ -z "$line" || "$line" == \#* ]] && continue
        echo "Bridge $line" >> $TORRC
```

```
    done < "$BRIDGE_FILE"
fi

# Debian tor package uses 'debian-tor' user
chown -R debian-tor:debian-tor /var/lib/tor /etc/tor

exec gosu debian-tor /usr/bin/tor -f $TORRC
```

27 launchd/com.nullexit.noise.plist

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>Label</key>
    <string>com.nullexit.noise</string>

    <!-- Same install-agnostic pattern as com.nullexit.wake-recovery: WorkingDirectory
         is __NULLEXIT_HOME__ (substitute your absolute install root at install time
         via sed), and the program arg is scripts/noise.sh RELATIVE to it, so
         BASH_SOURCE still resolves SCRIPT_DIR correctly. -->
    <key>WorkingDirectory</key>
    <string>__NULLEXIT_HOME__</string>

    <!-- `__boot`: governed SOLELY by NOISE_ENABLED in .env. It exits 0 cleanly when
         disabled (so KeepAlive never tight-loops) and runs the padding loop in the
         foreground when enabled. Loading this plist adds NO new option  it just means
         "auto-start cover traffic whenever NOISE_ENABLED=true". -->
    <key>ProgramArguments</key>
    <array>
        <string>/bin/bash</string>
        <string>scripts/noise.sh</string>
        <string>__boot</string>
    </array>

    <!-- Start immediately on launchctl load and re-run at every login/boot. -->
    <key>RunAtLoad</key>
    <true/>

    <!-- Same restart policy rationale as wake-recovery:
         SuccessfulExit=false  a clean exit 0 (the disabled-idle path) is NOT
                               relaunched, so a disabled job doesn't spin.
         Crashed=false         a SIGTERM from `launchctl bootout` is not treated as
                               a relaunch trigger, so stopping is clean and loop-free.
         Net effect: load once  padding runs at boot when enabled; bootout stops it. -->
    <key>KeepAlive</key>
    <dict>
        <key>SuccessfulExit</key>
        <false/>
        <key>Crashed</key>
        <false/>
    </dict>

    <!-- Cover traffic must not be suspended by App Nap when the user is idle. -->
    <key>ProcessType</key>
    <string>Background</string>

    <key>StandardOutPath</key>
    <string>/tmp/nullexit-noise.out.log</string>
    <key>StandardErrorPath</key>
    <string>/tmp/nullexit-noise.err.log</string>
</dict>
</plist>
```

28 launchd/com.nullexit.taildrop-watch.plist

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>Label</key>
    <string>com.nullexit.taildrop-watch</string>
```

```

    <!-- Same install-agnostic pattern as com.nullexit.wake-recovery /
         com.nullexit.noise: WorkingDirectory is __NULLEXIT_HOME__ (substituted
         with the absolute install root at install time), and the program arg
         is scripts/taildrop-watch.sh RELATIVE to it, so BASH_SOURCE inside the
         script still resolves correctly. -->
    <key>WorkingDirectory</key>
    <string>__NULLEXIT_HOME__</string>

    <key>ProgramArguments</key>
    <array>
        <string>/bin/bash</string>
        <string>scripts/taildrop-watch.sh</string>
    </array>

    <key>RunAtLoad</key>
    <true/>

    <!-- Unlike noise/wake-recovery there's no "disabled, sit idle" exit path
         here  `tailscale file get --loop` only returns on a real failure
         (tailscaled restarting, socket briefly gone), so always respawning is
         correct. launchd's built-in crash throttling covers a persistent
         failure without spinning. -->
    <key>KeepAlive</key>
    <true/>

    <key>ProcessType</key>
    <string>Background</string>

    <key>StandardOutPath</key>
    <string>/tmp/nullexit-taildrop.out.log</string>
    <key>StandardErrorPath</key>
    <string>/tmp/nullexit-taildrop.err.log</string>
</dict>
</plist>
```

29 launchd/com.nullexit.wake-recovery.plist

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>Label</key>
    <string>com.nullexit.wake-recovery</string>

    <!-- WorkingDirectory + relative program arg keeps this plist install-agnostic.
         The source-of-truth for the install path is __NULLEXIT_HOME__, which
         setup.sh or the user must substitute with their absolute nullexit
         install root at install time. Program arg is `scripts/watcher.sh`
         RELATIVE to WorkingDirectory, so launchd's evaluated cwd IS the
         install root  and BASH_SOURCE inside scripts/watcher.sh still
         resolves to the real <install>/scripts/watcher.sh, preserving the
         SCRIPT_DIR pattern. After install-time `sed -i '' 's|__NULLEXIT_HOME__|<abs_path>|'`,
         the plist is portable to any clone location. -->
    <key>WorkingDirectory</key>
    <string>__NULLEXIT_HOME__</string>

    <key>ProgramArguments</key>
    <array>
        <string>/bin/bash</string>
        <string>scripts/watcher.sh</string>
    </array>

    <!-- Start the daemon immediately on launchctl load (instead of waiting
         for the next user login). Combined with keepAlive below this means
         once installed: load once, runs forever across logouts/reboots. -->
    <key>RunAtLoad</key>
    <true/>

    <!-- Restart policy:
         SuccessfulExit=false   SIGTERM / clean exit does NOT relaunch (so
                                `launchctl bootout` cleanly stops, not loops).
         Crashed=false         HARD crashes ALSO do NOT relaunch. We've seen
                                that repeated sleep/wake cycles on macOS
                                Sonoma+ accumulate orphan watcher.sh PIDs
                                because SIGCONT/resume doesn't reliably kill
```

```
                                the prior instance. The watcher.sh script
                                now enforces single-instance via a manual
                                PID-file and process-identity check at the
                                top, so a re-launch on crash is
                                actively dangerous (it would skip the lock
                                and exit; better to surface the crash in
                                logs than to silently 0-process). -->
    <key>KeepAlive</key>
    <dict>
        <key>SuccessfulExit</key>
        <false/>
        <key>Crashed</key>
        <false/>
    </dict>

    <!-- Background hint to App Nap so it doesn't get suspended when
         the user is inactive. We're the entire reason this LaunchAgent
         needs to NOT nap. -->
    <key>ProcessType</key>
    <string>Background</string>

    <!-- launchd-level stdout/stderr (the bash launcher shell). The script
         itself writes to /tmp/nullexit-watcher.log via exec >>. These
         separate channels help us tell "launchd crashed" from "the bash
         wrapper piped to recover.sh errored". -->
    <key>StandardOutPath</key>
    <string>/tmp/nullexit-watcher.out.log</string>
    <key>StandardErrorPath</key>
    <string>/tmp/nullexit-watcher.err.log</string>
</dict>
</plist>
```

30 post-rules.txt

```
# NOTE: the REJECT rules below (DNS-over-TLS, QUIC) must be added BEFORE the blanket
# FORWARD ACCEPT rules. iptables stops at the first match, so if the ACCEPT rules ran
# first, all tailscale0 traffic (including port 853/443) would already be accepted and
# these REJECT rules would never be reached  they'd sit dead with a permanent 0-packet counter.

# Block DNS-over-TLS (Port 853) to force Android Private DNS to fallback to Port 53
iptables -A FORWARD -i tailscale0 -p tcp --dport 853 -j REJECT --reject-with tcp-reset
iptables -A FORWARD -i tailscale0 -p udp --dport 853 -j REJECT

# Block QUIC (UDP 443) to force Facebook/Google to fallback to TCP.
# UDP bypasses TCP MSS clamping, causing fragmentation. Over weak Wi-Fi (far from gateway),
# dropping a single fragment destroys the entire packet, stalling image downloads.
iptables -A FORWARD -i tailscale0 -p udp --dport 443 -j REJECT

iptables -A FORWARD -i tailscale0 -o tun0 -j ACCEPT
iptables -A FORWARD -i tun0 -o tailscale0 -j ACCEPT
iptables -A INPUT -i tailscale0 -j ACCEPT
iptables -A INPUT -i eth0 -p udp --dport 41642 -j ACCEPT
iptables -t nat -A POSTROUTING -o tun0 -j MASQUERADE
# Clamp MSS to 1120. With Tailscale overhead on top of WARP (each WireGuard encapsulation adds ~60 bytes),
# the safe MSS for double-tunneled traffic is 1120 to prevent strict-MSS endpoints from stalling.
# NOTE: this 1120 integer is rewritten at boot by toggle.sh from GATEWAY_MSS in .env
# (Gluetun's parser can't interpolate variables)  change GATEWAY_MSS in .env, not this line.
iptables -t mangle -A POSTROUTING -p tcp --tcp-flags SYN,RST SYN -j TCPMSS --set-mss 1120

iptables -t nat -A PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT --to-port 5335
iptables -t nat -A PREROUTING -i tailscale0 -p tcp --dport 53 -j REDIRECT --to-port 5335

# Also redirect DNS from Docker bridge interface so the host can reach AdGuard via localhost:53
iptables -t nat -A PREROUTING -i eth0 -p udp --dport 53 -j REDIRECT --to-port 5335
iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 53 -j REDIRECT --to-port 5335

# Block IPv6 exit-node traffic to prevent leakage (Gluetun/WARP is IPv4-only)
ip6tables -A FORWARD -i tailscale0 -j DROP
```

31 scripts/Dockerfile.routing-fix

```
FROM golang:1.22-alpine AS builder
WORKDIR /app
COPY logger/ ./logger/
RUN cd logger && go build -o logger main.go

FROM alpine:3.20
RUN apk add --no-cache iptables iproute2 curl ipset
COPY --from=builder /app/logger/logger /usr/local/bin/nullexit-logger
COPY routing-fix.sh /routing-fix.sh
CMD ["sh", "/routing-fix.sh"]
```

32 scripts/Dockerfile.rule-compiler

```
FROM golang:1.22-alpine AS builder
WORKDIR /app
COPY rule-compiler/ ./rule-compiler/
RUN cd rule-compiler && go build -o sync-rules main.go

FROM alpine:3.20
COPY --from=builder /app/rule-compiler/sync-rules /usr/local/bin/sync-rules
WORKDIR /app
CMD ["sync-rules"]
```

33 scripts/Toggle-Gateway.bat

```
<# :
@echo off
powershell -NoProfile -ExecutionPolicy Bypass -Command "Invoke-Expression ([System.IO.File]::ReadAllText('%~f0
    ') -replace '^(?s).*?#\s*>\s*\r?\n','')"
exit /b
#>
# Toggle-Gateway.bat (Hybrid Batch-PowerShell Toggler)  nullexit Gateway Toggler for Windows
# Mirrors the sophisticated error handling, checks, and automation of toggle.sh

#  1. Administrator Check & Elevation
$isAdmin = ([Security.Principal.WindowsPrincipal][Security.Principal.WindowsIdentity]::GetCurrent()).IsInRole([
    Security.Principal.WindowsBuiltInRole]::Administrator)
if (-not $isAdmin) {
    Write-Host "Requesting Administrator privileges..." -ForegroundColor Yellow
    Start-Process PowerShell -ArgumentList "-NoProfile -ExecutionPolicy Bypass -Command `"Invoke-Expression ([
     System.IO.File]::ReadAllText('$PSCommandPath') -replace '^(?s).*?#\s*>\s*\r?\n','')`"" -Verb RunAs
    Exit
}

# Set console output encoding to UTF8 for clean symbols
[Console]::OutputEncoding = [System.Text.Encoding]::UTF8
Clear-Host

Write-Host "" -ForegroundColor Green
Write-Host "           nullexit Windows Toggler           " -ForegroundColor Green
Write-Host "" -ForegroundColor Green
Write-Host ""

#  2. Helpers
function Reset-HostDNS {
    Write-Host "Restoring default DNS (DHCP/DHCP-assigned DNS)..." -ForegroundColor Cyan
    $activeAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
    foreach ($adapter in $activeAdapters) {
        Write-Host "  -> Resetting DNS on adapter: $($adapter.Name)" -ForegroundColor Gray
        Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ResetServerAddresses -ErrorAction
    SilentlyContinue
    }
}

function Hijack-HostDNS ($ip) {
    Write-Host "Hijacking DNS to route through AdGuard ($ip)..." -ForegroundColor Yellow
    $activeAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
    foreach ($adapter in $activeAdapters) {
        Write-Host "  -> Pointing DNS to $ip on adapter: $($adapter.Name)" -ForegroundColor Gray
        Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ServerAddresses $ip -ErrorAction
    SilentlyContinue
```

```
    }
}

function Get-HostTailscalePath {
    $paths = @(
        "$env:ProgramFiles\Tailscale\tailscale.exe",
        "$env:ProgramFiles (x86)\Tailscale\tailscale.exe",
        "tailscale.exe"
    )
    foreach ($path in $paths) {
        if (Test-Path $path) { return $path }
    }
    return $null
}

#  3. Detect State
# Prevent deadlocks by resetting DNS to default temporarily during checks
Reset-HostDNS

Write-Host "Checking Docker daemon status..." -ForegroundColor Cyan
& docker info >$null 2>&1
if ($LASTEXITCODE -ne 0) {
    Write-Host "Docker is not running. Starting Docker Desktop..." -ForegroundColor Yellow
    $dockerPath = "C:\Program Files\Docker\Docker\Docker Desktop.exe"
    if (Test-Path $dockerPath) {
        Start-Process $dockerPath
        Write-Host "Waiting for Docker daemon to become responsive..." -ForegroundColor Gray
        $timeout = 60
        while ($timeout -gt 0) {
            Start-Sleep -Seconds 2
            & docker info >$null 2>&1
            if ($LASTEXITCODE -eq 0) {
                Write-Host "Docker is ready!" -ForegroundColor Green
                break
            }
            $timeout -= 2
        }
        if ($timeout -le 0) {
            Write-Error "Docker failed to start in time. Aborting."
            Exit 1
        }
    } else {
        Write-Error "Docker Desktop executable not found at standard path. Please start Docker manually."
        Exit 1
    }
}

# Detect if gateway is already running
$runningWarp = docker compose ps --status running --format json | ConvertFrom-Json -ErrorAction
     SilentlyContinue | Where-Object { $_.Service -eq "warp" }
$isGatewayActive = ($null -ne $runningWarp)

#  4. Execution
if ($isGatewayActive) {
    #  STOP PATH
    Write-Host "Gateway is currently RUNNING. Disabling now..." -ForegroundColor Yellow
    Write-Host ""

    # Disconnect host Tailscale from exit node if installed
    $tsCli = Get-HostTailscalePath
    if ($null -ne $tsCli) {
        Write-Host "Disconnecting host Tailscale from exit node..." -ForegroundColor Cyan
        & $tsCli up --exit-node= --accept-dns=true >$null 2>&1
    }

    Write-Host "Stopping Docker containers..." -ForegroundColor Cyan
    docker compose down -t 5

    Reset-HostDNS
    Write-Host "Gateway has been successfully STOPPED." -ForegroundColor Green
} else {
    #  START PATH
    Write-Host "Gateway is currently STOPPED. Enabling now..." -ForegroundColor Yellow
    Write-Host ""

    # Clean up corrupted AdGuardHome configurations
    $confFile = "adguard\conf\AdGuardHome.yaml"
    if (Test-Path $confFile) {
        $fileSize = (Get-Item $confFile).Length
        if ($fileSize -eq 0) {
            Remove-Item $confFile -Force
```

```
            Write-Host "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop." -ForegroundColor Yellow
        }
    }

    Write-Host "Starting Docker containers..." -ForegroundColor Cyan
    docker compose up -d

    Write-Host "Waiting for container's Tailscale connection to offer Exit Node..." -ForegroundColor Cyan
    Write-Host "  (This can take up to 60 seconds...)" -ForegroundColor Gray

    $elapsed = 0
    $maxWait = 60
    $resolvedIP = $null

    while ($elapsed -lt $maxWait) {
        $statusJson = docker compose exec -T tailscale tailscale status --json 2>$null
        if ($null -ne $statusJson) {
            $status = $statusJson | ConvertFrom-Json -ErrorAction SilentlyContinue
            if ($null -ne $status -and $null -ne $status.Self) {
                # Check if it has a valid Tailscale IP and is running
                if ($status.Self.Online -and $status.Self.TailscaleIPs.Count -gt 0) {
                    $resolvedIP = $status.Self.TailscaleIPs[0]
                    break
                }
            }
        }
        Start-Sleep -Seconds 2
        $elapsed += 2
    }

    if ($null -ne $resolvedIP) {
        Write-Host "Tailscale IP resolved: $resolvedIP" -ForegroundColor Green
        Write-Host ""

        # Configure host Tailscale client to route through this exit node
        $tsCli = Get-HostTailscalePath
        if ($null -ne $tsCli) {
            Write-Host "Configuring host Tailscale to use Exit Node ($resolvedIP)..." -ForegroundColor Cyan
            & $tsCli up --exit-node=$resolvedIP --accept-dns=false >$null 2>&1
        }

        # Hijack DNS
        Hijack-HostDNS $resolvedIP
        Write-Host ""
        Write-Host "Gateway has been successfully STARTED." -ForegroundColor Green
    } else {
        Write-Host "Warning: Could not resolve gateway Tailscale IP. DNS hijacking skipped." -ForegroundColor Red
        # Fallback to .gateway_ip if it exists
        if (Test-Path ".gateway_ip") {
            $fallbackIP = Get-Content ".gateway_ip" -TotalCount 1
            if (-not [string]::IsNullOrWhiteSpace($fallbackIP)) {
                Write-Host "[Fallback] Using static IP from .gateway_ip: $fallbackIP" -ForegroundColor Yellow
                Hijack-HostDNS $fallbackIP
                Write-Host "Gateway STARTED (with static fallback IP)." -ForegroundColor Green
            }
        } else {
            Write-Host "Gateway started in offline/unauthenticated state. Run setup.sh if this is a fresh setup." -ForegroundColor Yellow
        }
    }
}

Write-Host ""
Write-Host "Press any key to exit..."
[void]$Host.UI.RawUI.ReadKey("NoEcho,IncludeKeyDown")
```

34 scripts/check_circuits.py

```
#!/usr/bin/env python3
import socket
import os
import re

def query_tor(port, command, password=""):
```

```
    try:
        s = socket.socket()
        s.settimeout(2)
        s.connect(('127.0.0.1', port))
        s.sendall(f"AUTHENTICATE \"{password}\"\r\n{command}\r\nQUIT\r\n".encode())
        response = b""
        while True:
            chunk = s.recv(4096)
            if not chunk:
                break
            response += chunk
        return response.decode()
    except Exception as e:
        return f"Error connecting to Tor Control Port: {e}"

def get_ports():
    ports = {}
    ports_file = "/tmp/nullexit-ports.env"
    if os.path.exists(ports_file):
        with open(ports_file, 'r') as f:
            for line in f:
                if line.startswith("export "):
                    # Split at first '=' to handle potential multiple '=' in value
                    parts = line.replace("export ", "").strip().split("=", 1)
                    if len(parts) == 2:
                        key, val = parts
                        try:
                            ports[key] = int(val)
                        except ValueError:
                            ports[key] = val
    return ports

def get_password(ports):
    # 1. Try ports.env
    password = ports.get("TOR_PASSWORD")
    if password:
        return password

    # 2. Try .env
    if os.path.exists(".env"):
        try:
            with open(".env", "r") as f:
                for line in f:
                    if line.startswith("ADGUARD_PASSWORD="):
                        return line.split("=", 1)[1].strip().strip('"\'')
                    if line.startswith("TOR_PASSWORD="):
                        return line.split("=", 1)[1].strip().strip('"\'')
        except Exception:
            pass

    # 3. Fallback
    return "nullexit"

def main():
    ports = get_ports()
    control_port = ports.get("TOR_CONTROL_PORT", 9051)
    password = get_password(ports)

    print(f"Connecting to Tor Control Port on 127.0.0.1:{control_port}...")
    status = query_tor(control_port, "GETINFO circuit-status", password)

    if "Error" in status:
        print(status)
        return

    print("\nActive Tor Circuits:")
    circuits = re.findall(r"(\d+)\s+BUILT\s+(.+?)\s+PURPOSE", status)
    if not circuits:
        print("No active built circuits found yet. Tor might still be bootstrapping.")
        return

    for circ_id, path_str in circuits:
        print(f"\nCircuit #{circ_id}:")
        nodes = path_str.split(",")
        for idx, node in enumerate(nodes):
            fingerprint = node.split("~")[0].replace("$", "")
            nickname = node.split("~")[1] if "~" in node else "unknown"

            # Fetch node IP
            ns_info = query_tor(control_port, f"GETINFO ns/id/{fingerprint}", password)
            ip = "unknown"
```

```
            for line in ns_info.split("\n"):
                if line.startswith("r "):
                    parts = line.split(" ")
                    if len(parts) >= 7:
                        ip = parts[6]
                        break

            # Fetch country
            country = "unknown"
            if ip != "unknown":
                country_info = query_tor(control_port, f"GETINFO ip-to-country/{ip}", password)
                for line in country_info.split("\n"):
                    if line.startswith("250-ip-to-country/"):
                        country = line.split("=")[1].strip().upper()
                        break

            role = "Guard" if idx == 0 else ("Exit" if idx == len(nodes) - 1 else "Middle")
            print(f"  [{role}] {nickname} ({ip}) - Country: {country}")

if __name__ == "__main__":
    main()
```

35 scripts/device-scramble.sh

```
#!/bin/bash
# scripts/device-scramble.sh  present a random, anonymous DEVICE identity on Wi-Fi.
#
# Randomizes the identifiers a network/observer uses to fingerprint your device:
#    hostname    ComputerName / LocalHostName / HostName (the DHCP hostname +
#                  the mDNS/Bonjour ".local" name broadcast on the LAN)
#    MAC (opt)   the layer-2 device address
#
# WHERE IT HELPS: open / no-auth Wi-Fi (caf, hotel, airport) where these IDs are
# the ONLY handle on you, and against co-located sniffers on any network. It does
# NOT hide you from a network you 802.1X-authenticate to  that login ties every
# session to your account regardless of MAC (see the FaceTime/threat-model notes).
#
# SAFETY: `scramble` saves your real identity first; `restore` puts it back.
# `test-mac` is FAIL-SAFE  it always restores your original MAC at the end, and
# detects Apple-Silicon spoof-blocking and MAC-gated networks without stranding you.
#
# Usage (privileged actions need sudo):
#   bash scripts/device-scramble.sh status            # show current + saved identity (no sudo)
#   sudo bash scripts/device-scramble.sh scramble     # random hostname now (add --mac to also rotate MAC)
#   sudo bash scripts/device-scramble.sh restore       # put your real identity back
#   sudo bash scripts/device-scramble.sh test-mac      # can this network take a random MAC? (auto-reverts)
#   bash scripts/device-scramble.sh --dry-run          # print the random values it would use (no sudo, no
    change)

set -uo pipefail
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
cd "$PROJECT_ROOT" || exit 1
IFACE=en0
SAVE="$PROJECT_ROOT/.device-identity.orig"   # gitignored
WARP_INHIBIT="/tmp/nullexit-warp-inhibit.marker"   # WARP Watcher skips shutdown while this exists

rand_mac() {  # locally-administered (0x02) + unicast (clear 0x01) first octet
  printf '%02x:%02x:%02x:%02x:%02x:%02x' \
    $(( (RANDOM & 0xFC) | 0x02 )) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM&255)) $((RANDOM
    &255))
}
rand_host() { printf 'MacBook-%04X' $((RANDOM & 0xFFFF)); }   # Apple-ish, blends in

cur_mac()  { ifconfig "$IFACE" ether 2>/dev/null | awk '/ether/{print $2}'; }
cur_ip()   { ipconfig getifaddr "$IFACE" 2>/dev/null; }
need_root(){ [ "$(id -u)" = 0 ] || { echo "device-scramble: this needs root  run with: sudo bash $0 $*" >&2;
    exit 1; }; }

cmd_status() {
  echo "current:"
  echo "  ComputerName  : $(scutil --get ComputerName 2>/dev/null)"
  echo "  LocalHostName : $(scutil --get LocalHostName 2>/dev/null)"
  echo "  HostName      : $(scutil --get HostName 2>/dev/null || echo '(unset)')"
  echo "  MAC ($IFACE)    : $(cur_mac)   IP: $(cur_ip)"
  if [ -f "$SAVE" ]; then echo; echo "saved original (restore target):"; sed 's/^/  /' "$SAVE"; fi
```

```
}

cmd_dry_run() {
  echo "dry-run  would set (no changes made, no sudo needed):"
  echo "  hostname  $(rand_host)"
  echo "  MAC       $(rand_mac)   (locally-administered, unicast)"
}

save_original() {
  [ -f "$SAVE" ] && return 0   # don't overwrite a real original with an already-scrambled one
  {
    echo "ComputerName=$(scutil --get ComputerName 2>/dev/null)"
    echo "LocalHostName=$(scutil --get LocalHostName 2>/dev/null)"
    echo "HostName=$(scutil --get HostName 2>/dev/null)"
    echo "MAC=$(cur_mac)"
  } > "$SAVE"
}

cmd_scramble() {
  need_root scramble
  save_original
  local h; h=$(rand_host)
  echo "device-scramble: hostname  $h"
  scutil --set ComputerName  "$h"
  scutil --set LocalHostName  "$h"
  scutil --set HostName       "$h"
  dscacheutil -flushcache 2>/dev/null || true
  if [ "${1:-}" = "--mac" ]; then
    local m; m=$(rand_mac)
    echo "device-scramble: MAC  $m (disassociate  set  reconnect)"
    networksetup -setairportpower "$IFACE" off; sleep 3   # can't set MAC while associated
    ifconfig "$IFACE" ether "$m" 2>/dev/null
    [ "$(cur_mac)" = "$m" ] || echo "  [!] MAC did not change  this driver refuses ifconfig spoofing on this
     hardware"
    networksetup -setairportpower "$IFACE" on
  fi
  echo "done. Restore your real identity with: sudo bash $0 restore"
}

cmd_restore() {
  need_root restore
  [ -f "$SAVE" ] || { echo "device-scramble: no saved identity at $SAVE  nothing to restore."; exit 1; }
  local cn lh hn mac
  cn=$(sed -n 's/^ComputerName=//p'  "$SAVE"); lh=$(sed -n 's/^LocalHostName=//p' "$SAVE")
  hn=$(sed -n 's/^HostName=//p' "$SAVE");      mac=$(sed -n 's/^MAC=//p' "$SAVE")
  [ -n "$cn" ] && scutil --set ComputerName  "$cn"
  [ -n "$lh" ] && scutil --set LocalHostName  "$lh"
  [ -n "$hn" ] && scutil --set HostName       "$hn"
  if [ -n "$mac" ] && [ "$(cur_mac)" != "$mac" ]; then
    networksetup -setairportpower "$IFACE" off; sleep 3   # disassociate before setting MAC
    ifconfig "$IFACE" ether "$mac" 2>/dev/null
    networksetup -setairportpower "$IFACE" on
  fi
  rm -f "$SAVE"
  echo "restored: ComputerName=$cn  MAC=$(cur_mac)"
}

# The safe experiment: does THIS network accept a random MAC? Always reverts.
# CRITICAL: macOS refuses to change the MAC while the interface is ASSOCIATED to
# an SSID, so we disassociate (Wi-Fi radio off) BEFORE setting it, then reconnect.
# (This is the fix for the naive first attempt that set it while connected.)
cmd_test_mac() {
  need_root test-mac
  # Protect the running gateway: inhibit the WARP Watcher so the brief Wi-Fi blip
  # isn't counted as tunnel death (which would auto-shutdown a healthy gateway).
  # Only manage the marker if we created it  don't clobber recover.sh's.
  local owned_inhibit=0
  if [ ! -f "$WARP_INHIBIT" ]; then touch "$WARP_INHIBIT" && owned_inhibit=1; fi
  trap '[ "$owned_inhibit" = 1 ] && rm -f "$WARP_INHIBIT"' EXIT INT TERM
  echo "gateway guard: WARP Watcher inhibited during the test (containers stay up regardless)."
  local orig; orig=$(cur_mac)
  [ -n "$orig" ] || { echo "no $IFACE MAC found"; exit 1; }
  echo "$orig" > /tmp/nullexit-mac.orig
  local rnd; rnd=$(rand_mac)
  echo "original MAC: $orig    trying random: $rnd"
  echo "step 1: disassociate (Wi-Fi off)  you can't set the MAC while connected"
  networksetup -setairportpower "$IFACE" off; sleep 3
  echo "step 2: set MAC while disassociated"
  ifconfig "$IFACE" ether "$rnd" 2>/dev/null
  local after; after=$(cur_mac)
  if [ "$after" != "$rnd" ]; then
```

```
    echo "RESULT: MAC still $after even when disassociated  this driver refuses ifconfig spoofing on neo."
    echo "        (Only Apple's native 'Private Wi-Fi Address' or a USB Wi-Fi adapter would work.)"
    networksetup -setairportpower "$IFACE" on; exit 0
  fi
  echo "   MAC CHANGED to $after with disassociate-first. Reconnecting to test the network"
  networksetup -setairportpower "$IFACE" on
  local ip=""; for _ in $(seq 1 15); do sleep 2; ip=$(cur_ip); [ -n "$ip" ] && break; done
  if [ -n "$ip" ]; then
    echo "RESULT:  SUCCESS  random MAC works AND got IP $ip. neo can spoof + this network accepts it."
  else
    echo "RESULT:   MAC spoof works, but no IP in ~30s  this network gates on your registered MAC (or 802.1X re
     -auth needed)."
  fi
  echo "restoring original MAC $orig "
  networksetup -setairportpower "$IFACE" off; sleep 3
  ifconfig "$IFACE" ether "$orig" 2>/dev/null
  networksetup -setairportpower "$IFACE" on
  ip=""; for _ in $(seq 1 15); do sleep 2; ip=$(cur_ip); [ -n "$ip" ] && break; done
  echo "restored: MAC=$(cur_mac)  IP=${ip:-<reconnecting>}"
}

case "${1:-status}" in
  status)    cmd_status ;;
  --dry-run) cmd_dry_run ;;
  scramble)  cmd_scramble "${2:-}" ;;
  restore)   cmd_restore ;;
  test-mac)  cmd_test_mac ;;
  --help|-h) sed -n '2,26p' "$0" ;;
  *) echo "Unknown: $1 (status|scramble [--mac]|restore|test-mac|--dry-run)" >&2; exit 2 ;;
esac
```

36 scripts/diagnose-host-leak.sh

```
#!/bin/bash
# scripts/diagnose-host-leak.sh  nullexit host-routing diagnostic
#
# Symptom this addresses: tests like https://www.whatismyip.com on the
# *HOST MAC* return the underlying physical ISP/ASN
# local ISP / campus network") even though the nullexit gateway is active and remote
# tailnet clients correctly egress via Cloudflare WARP. The container and
# the gateway itself are healthy  only the host's own routing is leaking.
#
# This script runs 7 checks in sequence, classifies the result into ONE
# of the three known leak scenarios, writes the full report to a
# timestamped file in the project root, and prints a verdict + a
# ready-to-run fix command on stdout. Pass --fix to apply the matching
# remediation in the same invocation and re-verify egress afterwards.
# See devref.md 10.30 for the deep-dive rationale.
#
# Usage:
#   bash scripts/diagnose-host-leak.sh              # diagnose + write report
#   bash scripts/diagnose-host-leak.sh --fix        # diagnose + fix + re-verify
#   bash scripts/diagnose-host-leak.sh --watch      # full baseline then loop checks every 60s
#   bash scripts/diagnose-host-leak.sh --watch 30   # same, with custom interval (seconds)
#   bash scripts/diagnose-host-leak.sh --help       # this message
#
# Multiple runs produce timestamped files (one per run) so historical
# reports can be diffed. The newest file is the most recent.

set -uo pipefail
# NOTE: errexit (`set -e`) is intentionally NOT enabled. Several checks
# below are EXPECTED to fail and that's diagnostic data (e.g. `tailscale
# status` returns non-zero when the daemon is wedged). Every command
# pipes to a helper that records the exit code without killing the script.

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
source "$SCRIPT_DIR/common.sh"
LOG_FILE="$PROJECT_ROOT/output.log"
REPORT_FILE="$PROJECT_ROOT/host-leak-diagnostic-$TIMESTAMP.txt"

#  Argument parsing
DO_FIX=false
DO_WATCH=false
WATCH_INTERVAL=60
```

```
WATCH_LOG="$PROJECT_ROOT/host-leak-watch.log"
case "${1:-}" in
  --fix) DO_FIX=true ;;
  --watch)
    DO_WATCH=true
    # Optional second argument: custom interval in seconds
    if [ -n "${2:-}" ] && [[ "$2" =~ ^[0-9]+$ ]]; then
      WATCH_INTERVAL="$2"
    fi
    ;;
  --help|-h)
    sed -n '2,31p' "$0"
    exit 0
    ;;
  "")
    : # default: diagnose only
    ;;
  *)
    echo "Unknown argument: $1" >&2
    echo "Run with --help for usage." >&2
    exit 2
    ;;
esac

#  Colours (auto-disabled when stdout is not a tty, e.g. piped)
if [ -t 1 ]; then
  RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
  BOLD='\033[1m'; NC='\033[0m'
else
  RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
fi

#  Output helpers  write to BOTH stdout and the report file
exec 3>> "$REPORT_FILE" || {
  printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
  exit 1
}
section() {
  local title="$1"
  printf '\n%b %s %b\n' "$BOLD" "$title" "$NC" >&3
  printf '\n%b %s %b\n' "$BOLD" "$title" "$NC"
}
line()   { printf '  %b\n' "$1" >&3; printf '  %b\n' "$1"; }
ok()     { line "${GREEN}${NC} $1"; }
warn()   { line "${YELLOW}${NC} $1"; }
fail()   { line "${RED}${NC} $1"; }
note()   { line "(no ${1})"; }

#  Watch-mode header (only shown when entering watch loop)
if [ "$DO_WATCH" = "true" ]; then
  echo ""
  echo ""
  echo " n u l l e x i t   w a t c h   m o d e"
  echo ""
  echo " Interval:        ${WATCH_INTERVAL}s"
  echo " Watch log:       $WATCH_LOG"
  echo ""
else
  echo ""
  echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c"
  echo ""
  echo " Timestamp (UTC): $TIMESTAMP"
  echo " Project root:    $PROJECT_ROOT"
  echo " Report file:     $REPORT_FILE"
  echo " Mode:            $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only (pass --
    fix to remediate)")"
  echo ""
  echo "" >&3
  echo "" >&3
  echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c" >&3
  echo "" >&3
  echo " Timestamp (UTC): $TIMESTAMP" >&3
  echo " Mode:            $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only")" >&3
fi

#  Always add Homebrew to PATH (same as toggle.sh / recover.sh)
export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"

#  Pure-bash timeout (macOS lacks GNU `timeout`)
run_with_timeout() {
  local timeout_sec="$1"
```

```
  shift
  if [ "$#" -eq 0 ]; then return 1; fi
  "$@" &
  local cmd_pid=$!
  (
    sleep "$timeout_sec"
    if kill -0 "$cmd_pid" 2>> "$LOG_FILE"; then
      kill -9 "$cmd_pid" 2>> "$LOG_FILE" || true
    fi
  ) >> "$LOG_FILE" 2>&1 &
  local watcher_pid=$!
  set +e
  wait "$cmd_pid"
  local exit_status=$?
  set -uo pipefail
  kill "$watcher_pid" 2>> "$LOG_FILE" || true
  return "$exit_status"
}

#
# Quick-check functions (used by --watch loop; return simple status strings)
# Each echoes its result so the caller can capture and compare.
#

quick_warp() {
  # Returns: "on", "off", "unknown"
  local out
  out=$(run_with_timeout 8 curl -s --max-time 6 \
        https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null || echo "")
  local warp
  warp=$(printf '%s' "$out" | awk -F'=' '/^warp=/{print $2; exit}')
  printf '%s' "${warp:-unknown}"
}

quick_ipv6() {
  # Returns: the IPv6 address if leaking, or empty string if clean
  local out
  out=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>/dev/null || echo "")
  if [ -n "$out" ] && printf '%s' "$out" | grep -qE '^[0-9a-fA-F:]+$'; then
    printf '%s' "$out"
  fi
}

quick_default_route() {
  # Returns: "utun" if default route is via Tailscale utun*,
  #          "wifi" if via physical Wi-Fi (192.168.x.x),
  #          "other" otherwise
  local route
  route=$(netstat -rn 2>/dev/null | awk '/^default/ {print $0; exit}')
  if [ -z "$route" ]; then
    printf '%s' "none"
  elif printf '%s' "$route" | grep -qE 'utun[0-9]+'; then
    printf '%s' "utun"
  elif printf '%s' "$route" | grep -qE '^(default|0\.0\.0\.0).*\b192\.168\.'; then
    printf '%s' "wifi"
  else
    printf '%s' "other"
  fi
}

#  Watch loop (runs after baseline diagnostic when --watch is passed)
run_watch_loop() {
  local baseline_warp="$1"
  local baseline_default="$2"
  local baseline_ipv6_leak="$3"
  local interval="$4"

  echo ""
  echo ""
  echo " Baseline established. Monitoring every ${interval}s..."
  echo "  warp:     $baseline_warp"
  echo "  default:  $baseline_default"
  echo ""

  #  Safety: warn if interval is very short (avoids hammering APIs)
  if [ "$interval" -lt 30 ]; then
    printf '%b  Short interval (%ss)  rate-limiting risk on external APIs%b\n' "$YELLOW" "$interval" "$NC"
    printf '  Consider 30s or longer for production use.\n'
  fi
  printf '[%s] WATCH START  warp=%s  default=%s  interval=%ss\n' \
    "$(date -u +%Y-%m-%dT%H:%M:%SZ)" "$baseline_warp" "$baseline_default" "$interval" \
```

```
    >> "$WATCH_LOG"

  local iteration=0
  local last_warp="$baseline_warp"
  local last_default="$baseline_default"
  local last_ipv6_ok=$([ "$baseline_ipv6_leak" = "false" ] && echo true || echo false)

  trap 'printf "\n%bWatch stopped by user. Log at: %s%b\n" "$YELLOW" "$WATCH_LOG" "$NC"; exit 0' INT TERM

  while true; do
    iteration=$((iteration + 1))

    local warp_now ipv6_now default_now
    warp_now=$(quick_warp)
    ipv6_now=$(quick_ipv6)
    default_now=$(quick_default_route)

    local ts
    ts=$(date -u +%H:%M:%S)

    #  WARP check
    if [ "$warp_now" != "$last_warp" ]; then
      if [ "$warp_now" = "off" ] || [ "$warp_now" = "unknown" ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bWARP FLIPPED: %s  %s%b  YOUR IP IS LEAKING!\n' \
          "$RED" "$ts" "$BOLD" "$last_warp" "$warp_now" "$NC"
        printf '[%s] ALERT  warp flipped  %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
      else
        printf '\n%b[%s]  warp recovered: %s  %s%b\n' \
          "$GREEN" "$ts" "$last_warp" "$warp_now" "$NC"
        printf '[%s] OK     warp recovered  %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
      fi
      last_warp="$warp_now"
    fi

    #  IPv6 check
    if [ -n "$ipv6_now" ]; then
      if [ "$last_ipv6_ok" = true ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bIPv6 LEAK DETECTED  %s%b\n' \
          "$RED" "$ts" "$BOLD" "$ipv6_now" "$NC"
        printf '[%s] ALERT  IPv6 leak  %s\n' "$ts" "$ipv6_now" >> "$WATCH_LOG"
        last_ipv6_ok=false
      fi
    else
      if [ "$last_ipv6_ok" = false ]; then
        printf '%b[%s]  IPv6 leak resolved%b\n' "$GREEN" "$ts" "$NC"
        printf '[%s] OK     IPv6 resolved\n' "$ts" >> "$WATCH_LOG"
      fi
      last_ipv6_ok=true
    fi

    #  Default route check
    if [ "$default_now" != "$last_default" ]; then
      if [ "$default_now" != "utun" ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bDEFAULT ROUTE CHANGED: %s  %s%b  traffic may bypass WARP!\n' \
          "$RED" "$ts" "$BOLD" "$last_default" "$default_now" "$NC"
        printf '[%s] ALERT  route changed  %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
      else
        printf '%b[%s]  default route recovered: %s  %s%b\n' \
          "$GREEN" "$ts" "$last_default" "$default_now" "$NC"
        printf '[%s] OK     route recovered  %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
      fi
      last_default="$default_now"
    fi

    #  Heartbeat (every 10th iteration or if all OK)
    if [ $((iteration % 10)) -eq 0 ]; then
      printf '[%s] #%d  warp=%s  ipv6=%s  route=%s\n' \
        "$ts" "$iteration" "$warp_now" "${ipv6_now:-none}" "$default_now" >> "$WATCH_LOG"
      printf '\r%b[%s] #%d  warp=%s  route=%s  (Ctrl+C to stop)%b' \
        "$GREEN" "$ts" "$iteration" "$warp_now" "$default_now" "$NC"
    fi

    sleep "$interval"
  done
}

#
# Begin main diagnostic (skipped in non-watch modes if we're just watching)
```

```
#

ACTIVE_SERVICE=$(get_active_service)
ENO_SERVICE=$(networksetup -listnetworkserviceorder 2>> "$LOG_FILE" \
              | grep -B 1 "Device: en0" | head -1 \
              | sed -E 's/^\([0-9\*]+\) //' || true)
[ -z "$ENO_SERVICE" ] && ENO_SERVICE="Wi-Fi"

# Resolve the gateway's Tailscale IP.
GATEWAY_TS_IP=""
if [ -f "$PROJECT_ROOT/.gateway_ip" ]; then
  GATEWAY_TS_IP=$(cat "$PROJECT_ROOT/.gateway_ip" 2>> "$LOG_FILE" \
                  | tr -d '\r' | awk 'NR==1{print $1; exit}')
fi
if [ -z "$GATEWAY_TS_IP" ] && command -v docker >> "$LOG_FILE" 2>&1 \
   && docker compose ps --status running 2>/dev/null | grep -q 'tailscale'; then
  GATEWAY_TS_IP=$(run_with_timeout 10 docker compose exec -T tailscale \
                  tailscale ip -4 2>> "$LOG_FILE" \
                  | tr -d '\r' | awk 'NR==1{print $1; exit}')
fi

#  Scenario classification state
SCENARIO=""
WARP_STATUS=""
SOCKS5_ENABLED=false
TS_ON_MESH=false
DEFAULT_VIA_UTUN=false
IPV6_LEAK=false

#
section "1/8  Host Tailscale mesh status"
#
if ! command -v tailscale >> "$LOG_FILE" 2>&1; then
  fail "tailscale CLI not on PATH  install via 'brew install tailscale' then re-run"
  echo "  (a Tailscale daemon is required for the exit-node path to work)"
else
  set +e
  TS_OUT=$(run_with_timeout 5 tailscale status 2>&1)
  TS_EXIT=$?
  set -uo pipefail
  TS_HEAD=$(printf '%s\n' "$TS_OUT" | head -5 | tr -d '\r')
  if [ "$TS_EXIT" -eq 137 ]; then
    fail "tailscaled daemon is wedged/unresponsive (5s timeout)"
    echo "  Fix path: 'sudo brew services restart tailscale'"
  elif printf '%s' "$TS_OUT" | grep -qE 'Logged out|not logged in|expired'; then
    fail "tailscale is logged out / auth expired on this host"
    echo "  Fix path: 'sudo tailscale up' (browser auth once)"
  elif printf '%s' "$TS_OUT" | grep -q '100\.'; then
    ok "Host is on tailnet (100.x.x.x visible in status)"
    TS_ON_MESH=true
    if printf '%s' "$TS_OUT" | grep -q "$GATEWAY_TS_IP"; then
      ok "Gateway $GATEWAY_TS_IP visible in host's peer list"
    elif [ -n "$GATEWAY_TS_IP" ]; then
      warn "Gateway $GATEWAY_TS_IP NOT in host's peer list  possible auth/key mismatch"
    fi
    line " first 5 lines of 'tailscale status' "
    line "$(printf '%s' "$TS_HEAD" | awk '{print "    "$0}')"
  else
    fail "tailscale output unrecognized:"
    line "$(printf '%s' "$TS_OUT" | head -3 | awk '{print "    "$0}')"
  fi
fi

#
section "2/8  Active SOCKS5 proxy on Wi-Fi"
#
SOCKS_OUT=$(networksetup -getsocksfirewallproxy "$ACTIVE_SERVICE" 2>> "$LOG_FILE" \
            || networksetup -getsocksfirewallproxy Wi-Fi 2>> "$LOG_FILE" \
            || echo "Could not query SOCKS5 proxy state")
if printf '%s' "$SOCKS_OUT" | grep -q "Enabled: Yes"; then
  fail "SOCKS5 proxy IS enabled on $ACTIVE_SERVICE  toggle.sh fell back to SOCKS5 mode last run"
  SOCKS5_ENABLED=true
  line "  Full state:"
  line "$(printf '%s' "$SOCKS_OUT" | awk '{print "    "$0}')"
  line "  Browse on macOS ignores system SOCKS5 by default  traffic bypasses WARP."
else
  ok "SOCKS5 proxy is OFF on $ACTIVE_SERVICE  exit-node path should be active"
fi

#
section "3/8  Egress IP + WARP tunnel verification (definitive test)"
```

```
#
CDN_OUT=$(run_with_timeout 8 curl -s --max-time 6 \
          https://www.cloudflare.com/cdn-cgi/trace 2>&1 || echo "(curl failed)")
CDN_WARP=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^warp=/{print $2; exit}')
CDN_IP=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^ip=/{print $2; exit}')
CDN_COLO=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^colo=/{print $2; exit}')
WARP_STATUS="$CDN_WARP"
if [ "$CDN_WARP" = "on" ]; then
  ok "warp=on  tunnel is hot, host traffic IS going through Cloudflare WARP"
  line "  Public IP:   $CDN_IP"
  line "  Cloudflare:   ${CDN_COLO:-unknown colo}"
  line "   whatismyip.com's 'local ISP' result is its ISP-database error, not yours."
elif [ "$CDN_WARP" = "off" ]; then
  fail "warp=off  host traffic is NOT going through WARP"
  line "  Public IP:   $CDN_IP  (NOT Cloudflare)"
  warn "This is the actual leak."
else
  warn "could not determine warp status (curl failed or returned unexpected shape)"
  line "  Raw output: $(printf '%s' "$CDN_OUT" | head -3 | tr '\n' '|')"
fi

#
section "4/8  IPv6 leak probe (bypasses IPv4 exit-node on campus networks)"
#
IPV6_OUT=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
if [ -n "$IPV6_OUT" ] && printf '%s' "$IPV6_OUT" | grep -qE '^[0-9a-fA-F:]+$'; then
  fail "host has live IPv6 egress  $IPV6_OUT (A6 record bypasses WARP)"
  IPV6_LEAK=true
  line "  Tailscale exit-node advertisement is IPv4-only, so IPv6 traffic skips it."
  line "  Quick-fix: 'sudo networksetup -setv6off $ACTIVE_SERVICE'"
else
  ok "no IPv6 egress detected (good  IPv6 won't bypass the tunnel)"
fi

#
section "5/8  Host egress route for real traffic (route -n get 1.1.1.1)"
#
# WHY NOT 'netstat -rn | grep default':
# On macOS, Tailscale does NOT replace the physical default route installed by
# DHCP. Instead it adds a second interface-scoped route bound to the utun
# interface that the kernel prefers for real traffic forwarding. The literal
# "default" entry (192.168.x.x via en0) is left untouched in the table as a
# BSD routing quirk  it's not lying, it's just answering a different question.
# The correct question is: "where does a real destination actually resolve?"
# `route -n get 1.1.1.1` asks the kernel's forwarding logic directly.
ROUTE_GET=$(route -n get 1.1.1.1 2>/dev/null)
ROUTE_IFACE=$(echo "$ROUTE_GET" | awk '/interface:/{print $2}')
ROUTE_GATEWAY=$(echo "$ROUTE_GET" | awk '/gateway:/{print $2}')
if [ -z "$ROUTE_IFACE" ]; then
  warn "route -n get 1.1.1.1 returned no interface  routing table may be empty"
else
  line "  interface: $ROUTE_IFACE  gateway: ${ROUTE_GATEWAY:-n/a}"
  if echo "$ROUTE_IFACE" | grep -qE '^utun[0-9]+'; then
    ok "1.1.1.1 resolves via $ROUTE_IFACE  Tailscale interface-scoped route is winning"
    DEFAULT_VIA_UTUN=true
  elif echo "$ROUTE_IFACE" | grep -qE '^en[0-9]+'; then
    fail "1.1.1.1 resolves via $ROUTE_IFACE (physical Wi-Fi)  exit-node interface route not active"
  else
    warn "1.1.1.1 resolves via unexpected interface: $ROUTE_IFACE"
  fi
fi

#
section "6/8  Last toggle.sh START-relevant lines in output.log (if any)"
#
if [ ! -f "$LOG_FILE" ]; then
  note "output.log"
  line "  toggle.sh was never run from this directory. Run './toggle.sh' first,"
  line "  then re-run this diagnostic."
else
  HITS=$(grep -E 'Exit node enabled|SKIP_EXIT_NODE|Pre-flight|Falling back to SOCKS|Starting|Successfully' \
         "$LOG_FILE" 2>/dev/null | tail -8 || true)
  if [ -z "$HITS" ]; then
    note "matching lines in output.log"
  else
    line "$(printf '%s\n' "$HITS" | awk '{print "    "$0}')"
  fi
fi

#
section "7/8  Gateway IP we should be pointing at"
```

```
#
if [ -n "$GATEWAY_TS_IP" ]; then
  ok "gateway Tailscale IP: $GATEWAY_TS_IP"
else
  warn "could not resolve gateway Tailscale IP"
  line "  (looked in .gateway_ip and tried docker compose exec tailscale ip -4)"
fi

#
section "8/8  Tailscale System Extension (App Store conflict check)"
#
SE_PENDING_UNINSTALL=false
SE_ACTIVE=false
if command -v systemextensionsctl >/dev/null 2>&1; then
  SE_OUT=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale')
  if [ -n "$SE_OUT" ]; then
    line "  Tailscale System Extension(s) detected:"
    line "$(printf '%s' "$SE_OUT" | awk '{print "    "$0}')"
    if printf '%s' "$SE_OUT" | grep -q 'waiting to uninstall'; then
      SE_PENDING_UNINSTALL=true
      warn "System Extension is pending uninstall on next reboot."
    else
      SE_ACTIVE=true
      warn "Active Tailscale System Extension detected. This will conflict with Homebrew Tailscale."
    fi
  else
    ok "no Tailscale System Extension detected"
  fi
else
  note "systemextensionsctl not available (not on macOS or permission restricted)"
fi

#
section "CLASSIFICATION    applies ONE of the four known scenarios"
#

if [ "$WARP_STATUS" = "on" ] && [ "$IPV6_LEAK" = "false" ]; then
  SCENARIO="OK"
elif [ "$SE_PENDING_UNINSTALL" = "true" ]; then
  SCENARIO="SE_PENDING"
elif [ "$SOCKS5_ENABLED" = "true" ]; then
  SCENARIO="A"
elif [ "$IPV6_LEAK" = "true" ]; then
  SCENARIO="B"
elif [ "$DEFAULT_VIA_UTUN" = "false" ] && [ "$TS_ON_MESH" = "true" ]; then
  SCENARIO="C"
else
  SCENARIO="UNKNOWN"
fi

case "$SCENARIO" in
  SE_PENDING)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario SE_PENDING  Tailscale System Extension Pending Uninstall${NC}"
    echo "  A prior App Store Tailscale System Extension is pending uninstall."
    echo "  Since System Integrity Protection (SIP) is enabled, macOS blocks manual"
    echo "  uninstallation of terminated system extensions until the next reboot."
    echo "  Brew-managed tailscaled cannot engage the Network Extension slot until this is resolved."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo shutdown -r now"
    echo "    # After reboot, run:  ./toggle.sh"
    ;;
  A)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario A  SOCKS5 fallback lane${NC}"
    echo "  toggle.sh's pre-flight checks failed last run. The script activated"
    echo "  SOCKS5 + local DNS proxy as a fallback. DNS gets hijacked but"
    echo "  browsers on macOS ignore the system SOCKS5  traffic egresses"
    echo "  direct through the local network."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    echo 'GATEWAY_BYPASS_PING=true' >> $PROJECT_ROOT/.env"
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
    if [ -n "$GATEWAY_TS_IP" ]; then
    echo "        --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
    else
    echo "        --exit-node=\$GATEWAY_TS_IP --exit-node-allow-lan-access=true  # fill in $GATEWAY_TS_IP"
    fi
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
    echo "    ./toggle.sh"
```

```
    ;;
  B)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario B  IPv6 leak over dual-stack campus/local IPv6${NC}"
    echo "  Tailscale exit-node advertises only IPv4. macOS sends AAAA queries"
    echo "  straight out en0, which is dual-stack on most campus APs  IPv6"
    echo "  traffic bypasses the WARP tunnel entirely."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo networksetup -setv6off $ACTIVE_SERVICE"
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
    echo "    # Re-enable IPv6 later with:  sudo networksetup -setv6automatic $ACTIVE_SERVICE"
    ;;
  C)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario C  Tailscale route-freeze${NC}"
    echo "  Host's default route points at the Wi-Fi gateway instead of the"
    echo "  Tailscale utun* interface. The exit-node preference is set but the"
    echo "  routing-table assertion did not take effect (devref 10.26)."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo brew services restart tailscale"
    echo "    sudo route delete -net 0.0.0.0/1"
    echo "    sudo route delete -net 128.0.0.0/1"
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
    if [ -n "$GATEWAY_TS_IP" ]; then
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
    echo "        --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
    else
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true --exit-node=\$TS_IP
     --exit-node-allow-lan-access=true"
    fi
    echo "    ./toggle.sh"
    ;;
  OK)
    echo ""
    echo -e "  ${GREEN}${BOLD}VERDICT:  No host leak detected  tunnel is working${NC}"
    echo "  The local ISP string from whatismyip.com is its IP-database"
    echo "  misclassifying the Cloudflare egress IP.  CDN-CGI's trace endpoint"
    echo "  (which is what we use here) reports warp=on because that endpoint"
    echo "  is hosted by Cloudflare itself and can see its own edge."
    echo "  If you want a third-party confirmation:"
    echo "    curl -s https://api.ipify.org; echo"
    ;;
  *)
    echo ""
    echo -e "  ${YELLOW}${BOLD}VERDICT:  Could not classify automatically${NC}"
    echo "  The diagnostic data did not match any known scenario cleanly."
    echo "  Likely causes:"
    echo "    - tailscaled is logged out / expired (see Section 1)"
    echo "    - Docker is not running (gateway containers down)"
    echo "    - .gateway_ip missing AND docker compose exec failed"
    echo "  Inspect the report at: $REPORT_FILE"
    ;;
esac
echo ""

#
# Write the verdict block to the report file
#
SECOND_BLOCK=$(cat <<EOF


 C L A S S I F I C A T I O N   R E S U L T

 Scenario: $SCENARIO
 WARP egress:  $WARP_STATUS  (cdn-cgi/trace)
 SOCKS5 lane:  $([ "$SOCKS5_ENABLED" = true ] && echo "ENABLED (browsers bypass)" || echo "off")
 IPv6 leak:    $([ "$IPV6_LEAK" = true ] && echo "YES (campus IPv6 egress)" || echo "no")
 Default via:  $([ "$DEFAULT_VIA_UTUN" = true ] && echo "utun* (Tailscale)" || echo "Wi-Fi gateway")
 Host on mesh: $([ "$TS_ON_MESH" = true ] && echo "yes" || echo "no")
 SE pending:   $([ "$SE_PENDING_UNINSTALL" = true ] && echo "yes" || echo "no")
 Gateway IP:   ${GATEWAY_TS_IP:-unresolved}
EOF
)
printf '%s\n' "$SECOND_BLOCK" >&3

#
# --fix mode
#
if [ "$DO_FIX" = "true" ] && [ "$SCENARIO" != "OK" ] && [ "$SCENARIO" != "UNKNOWN" ]; then
```

```
  echo ""
  echo ""
  echo " A P P L Y I N G   F I X  (scenario $SCENARIO)"
  echo ""
  echo ""

  case "$SCENARIO" in
    SE_PENDING)
      warn "System Extension uninstall is pending on next reboot."
      line "Please run 'sudo shutdown -r now' to reboot the system."
      ;;
    A)
      if [ -n "$GATEWAY_TS_IP" ]; then
        line " writing GATEWAY_BYPASS_PING=true to .env"
        ENV="$PROJECT_ROOT/.env"
        [ -f "$ENV" ] || touch "$ENV"
        if grep -q '^GATEWAY_BYPASS_PING=' "$ENV" 2>/dev/null; then
          sed -i '' 's/^GATEWAY_BYPASS_PING=.*/GATEWAY_BYPASS_PING=true/' "$ENV"
        else
          printf '\nGATEWAY_BYPASS_PING=true\n' >> "$ENV"
        fi
        line " re-asserting tailscale exit-node (with --reset)"
        sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
             --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
             >> "$LOG_FILE" 2>&1 \
          && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
          || warn "tailscale up did not return success (check $LOG_FILE)"
      else
        warn "no gateway Tailscale IP resolved  skipping tailscale up"
        warn "fix manually: 'sudo tailscale up --reset ... --exit-node=<TS_IP>'"
      fi
      ;;
    B)
      line " disabling IPv6 on $ACTIVE_SERVICE while gateway is up"
      sudo -n networksetup -setv6off "$ACTIVE_SERVICE" >> "$LOG_FILE" 2>&1 \
        && ok "IPv6 disabled on $ACTIVE_SERVICE (re-enable with 'setv6automatic')" \
        || warn "networksetup -setv6off failed (check $LOG_FILE)"
      ;;
    C)
      line " restarting tailscaled (route-table fix)"
      run_with_timeout 15 brew services restart tailscale >> "$LOG_FILE" 2>&1 \
        || run_with_timeout 15 sudo -n brew services restart tailscale >> "$LOG_FILE" 2>&1
      sleep 3
      line " flushing stale host routes + DNS cache"
      sudo -n route delete -net 0.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
      sudo -n route delete -net 128.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
      sudo -n dscacheutil -flushcache >> "$LOG_FILE" 2>&1 || true
      sudo -n killall -HUP mDNSResponder >> "$LOG_FILE" 2>&1 || true
      if [ -n "$GATEWAY_TS_IP" ]; then
        line " re-asserting exit-node after restart"
        sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
             --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
             >> "$LOG_FILE" 2>&1 \
          && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
          || warn "tailscale up did not return success (check $LOG_FILE)"
      fi
      ;;
  esac

  echo ""
  echo "Re-verifying after fix..."
  sleep 3
  CDN_WARP_AFTER=$(run_with_timeout 8 curl -s --max-time 6 \
      https://www.cloudflare.com/cdn-cgi/trace 2>&1 \
      | awk -F'=' '/^warp=/{print $2; exit}')
  IPV6_AFTER=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
  printf '\n[AFTER FIX]\n' >&3
  printf '  warp status:    %s\n' "$CDN_WARP_AFTER" >&3
  printf '  IPv6 egress:     %s\n' "${IPV6_AFTER:-none}" >&3
  line "  warp status:    ${CDN_WARP_AFTER:-unknown}"
  line "  IPv6 egress:    ${IPV6_AFTER:-none}"
  if [ "$CDN_WARP_AFTER" = "on" ] && [ -z "$IPV6_AFTER" ]; then
    echo ""
    ok "Fix verified  host is now routing through WARP. Remote clients unaffected."
  else
    echo ""
    warn "Fix did not fully take effect on first try. Likely causes:"
    line "   tailscaled needed a longer settling window (re-run diagnostic in 30s)"
    line "   The classified scenario was wrong; paste this report for triage."
    line "   Tailscale needs a full auth refresh: 'sudo tailscale up'"
  fi
```

```
fi

echo ""
echo ""
echo "  Full report written to:"
echo "    $REPORT_FILE"
echo ""
echo ""

#
# --watch mode: enter the continuous monitoring loop after baseline diagnostic
#
if [ "$DO_WATCH" = "true" ]; then
  # Build baseline from the full diagnostic just completed
  BASELINE_WARP="$WARP_STATUS"
  BASELINE_DEFAULT=$([ "$DEFAULT_VIA_UTUN" = "true" ] && echo "utun" || echo "wifi/other")

  if [ "$SCENARIO" != "OK" ]; then
    warn "Baseline diagnostic found issues (scenario: $SCENARIO)."
    line "  Watch loop will still run  alerts fire on any STATE CHANGE from current baseline."
    line "  Fix the issue first? Re-run with:  bash scripts/diagnose-host-leak.sh --fix"
    echo ""
  fi

  run_watch_loop "$BASELINE_WARP" "$BASELINE_DEFAULT" "$IPV6_LEAK" "$WATCH_INTERVAL"
fi

exit 0
```

37 scripts/dns-proxy.py

```
#!/usr/bin/env python3
"""
dns-proxy.py  Local DNS proxy for nullexit gateway.

Listens on UDP port 53, receives DNS queries, forwards them over TCP
to AdGuard at 127.0.0.1:5354 (the Docker port mapping), and returns
the response via UDP. Handles the DNS-over-TCP wire format (2-byte
length prefix) correctly  unlike socat which just forwards raw bytes.
"""

import socket
import struct
import sys
import os
from concurrent.futures import ThreadPoolExecutor

LISTEN_ADDR = os.environ.get("LISTEN_ADDR", "127.0.0.1")
LISTEN_PORT = int(os.environ.get("LISTEN_PORT", 53))
TARGET_HOST = os.environ.get("TARGET_HOST", "127.0.0.1")
TARGET_PORT = int(os.environ.get("TARGET_PORT", 5354))

# When invoked via `sudo -n`, custom env vars set by the caller are stripped by
# sudo's env_reset unless allowlisted in sudoers  which the shipped NOPASSWD
# rule for this exact script (no bash -c wrapper, no env prefix) does not do.
# A caller that needs non-default values drops them here *before* sudo'ing the
# literal whitelisted command; this process picks them up on top of the
# env-var/default values above. Same bounded local-TOCTOU class already
# accepted for /tmp/nullexit-ports.env (see devref.md 15.9.5)  restrictive
# perms, not a hard security boundary against a co-resident attacker.
_CONF = "/tmp/nullexit-dns-proxy.conf"
if os.path.isfile(_CONF):
    with open(_CONF) as _f:
        for _line in _f:
            _line = _line.strip()
            if not _line or _line.startswith("#") or "=" not in _line:
                continue
            _k, _v = _line.split("=", 1)
            if _k == "LISTEN_ADDR":
                LISTEN_ADDR = _v
            elif _k == "LISTEN_PORT":
                LISTEN_PORT = int(_v)
            elif _k == "TARGET_HOST":
                TARGET_HOST = _v
            elif _k == "TARGET_PORT":
                TARGET_PORT = int(_v)
```

```

def handle_query(data: bytes, client_addr: tuple, sock: socket.socket) -> None:
    """Forward a single DNS query over TCP and send the response back via UDP."""
    try:
        tcp = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        tcp.settimeout(5)
        tcp.connect((TARGET_HOST, TARGET_PORT))

        # DNS over TCP: prepend 2-byte length (big-endian)
        tcp.sendall(struct.pack("!H", len(data)) + data)

        # Read the 2-byte response length
        raw_len = tcp.recv(2)
        if len(raw_len) < 2:
            tcp.close()
            return
        resp_len = struct.unpack("!H", raw_len)[0]

        # Read the full response
        response = b""
        while len(response) < resp_len:
            chunk = tcp.recv(resp_len - len(response))
            if not chunk:
                break
            response += chunk

        tcp.close()

        # Send the response back over UDP (strip TCP length prefix)
        if response:
            sock.sendto(response, client_addr)
    except Exception:
        # Silently drop failed queries (malformed, timeout, etc.)
        pass


def main() -> None:
    # Create UDP socket
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
    sock.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)

    try:
        sock.bind((LISTEN_ADDR, LISTEN_PORT))
    except PermissionError:
        print(f"dns-proxy: ERROR  need root to bind port {LISTEN_PORT}", flush=True)
        sys.exit(1)
    except OSError as e:
        print(f"dns-proxy: ERROR  {e}", flush=True)
        sys.exit(1)

    print(f"dns-proxy: listening on UDP {LISTEN_ADDR}:{LISTEN_PORT}  TCP {TARGET_HOST}:{TARGET_PORT}", flush=
     True)

    # Use a thread pool to handle concurrent DNS queries efficiently without thread exhaustion
    executor = ThreadPoolExecutor(max_workers=64)

    while True:
        try:
            data, client_addr = sock.recvfrom(4096)  # Support EDNS0 (up to 4096 bytes) to prevent truncation
            executor.submit(handle_query, data, client_addr, sock)
        except KeyboardInterrupt:
            break
        except Exception:
            pass

    executor.shutdown(wait=False)
    sock.close()


if __name__ == "__main__":
    main()
```

38 scripts/dns_anomaly_detector.py

```
#!/usr/bin/env python3
```

```
"""
nullexit DNS anomaly detector (v2)  report-only DNS-tunneling / DGA / C2 finder.

Replaces the v1 whole-FQDN Shannon-entropy + mean3 demo (which measured the
wrong scope, was fooled by short encoded labels, assumed a normal distribution
DNS names don't have, and ignored the volumetric signal that actually defines
tunneling). v2 keeps the "statistical, data-based, 100% local, no third-party
deps" spirit but detects far better:

  1. n-gram (character bigram) improbability of the SUBDOMAIN labels only,
     trained on YOUR querylog. Robust on short strings where Shannon entropy
     fails (entropy is capped at log2(len)); scores "doesn't look like real
     naming", which is what encoded tunneling labels are.
  2. per-registered-domain AGGREGATION (unique-subdomain count, query volume,
     suspicious record types)  a stream of high-improbability names under ONE
     domain is the real tunneling tell, not any single name.
  3. robust thresholds (median + kMAD, resistant to CDN/cloud hash outliers)
     instead of mean3 on a non-normal distribution.
  4. a data-driven allowlist of YOUR frequent domains, so CloudFront/S3/telemetry
     hashes don't false-positive.

Report-only by design: a privacy gateway must never silently drop DNS on a
statistical trip (a false positive = "my internet broke"). It surfaces findings;
promoting a domain to an actual block stays a human/AdGuard-blocklist decision.

Footprint: single streaming pass (never loads the log into RAM); the model is a
~4040 bigram table + a few floats + a small allowlist (a few KB on disk); per-
domain state is a handful of ints plus a subdomain set capped at 64.

Exit codes (scan): 0 = ran, clean  1 = ran, anomalies found  2 = FAILED to run
(no/corrupt model, unreadable or empty querylog)  a failure is never reported as
"clean". selftest exits 3 if the detection logic itself is broken.

Usage:
  dns_anomaly_detector.py learn [--log PATH] [--out PATH]   # train from querylog
  dns_anomaly_detector.py scan  [--log PATH] [--model PATH] [--recent N] [--quiet]
  dns_anomaly_detector.py selftest                          # prove it detects (exit 3 if broken)
  dns_anomaly_detector.py demo                              # human-readable showcase
"""
import argparse
import json
import math
import os
import random
import re
import sys
import time
from collections import defaultdict

DEFAULT_LOG = os.path.join(os.path.dirname(os.path.abspath(__file__)),
                           "..", "adguard", "work", "data", "querylog.json")
DEFAULT_MODEL = os.path.join(os.path.dirname(os.path.abspath(__file__)),
                             "..", ".dns_baseline.json")
OUTPUT_LOG = os.path.join(os.path.dirname(os.path.abspath(__file__)), "..", "output.log")


def log_fail(msg):
    """Persist a hard failure to output.log (besides stderr) so a broken detector
    is on the record even when run headless (launchd/sweep/CI). Heisenbug-safe by
    construction: it only ever APPENDS to a separate file  it touches neither
    stdout/stderr (which sweep captures) nor any detection state, and a logging
    error is swallowed so it can never change what the detector does or returns."""
    try:
        with open(OUTPUT_LOG, "a", encoding="utf-8") as f:
            f.write(f"[{time.strftime('%Y-%m-%d %H:%M:%S')}] [dns-anomaly] FAIL: {msg}\n")
    except OSError:
        pass

ALPHABET = "abcdefghijklmnopqrstuvwxyz0123456789-_"
AIDX = {c: i for i, c in enumerate(ALPHABET)}
OTHER = len(ALPHABET)          # bucket for any char outside the alphabet
NSYM = len(ALPHABET) + 1
BEG = NSYM                     # start-of-label marker
NSTATES = NSYM + 1

# Minimal two-level public suffixes so eTLD+1 grouping is stable without shipping
# the full ~200 KB Public Suffix List. Exact PSL isn't needed  we only need a
# consistent grouping key per registered domain.
TWO_LEVEL = {
    "co.uk", "org.uk", "ac.uk", "gov.uk", "co.jp", "co.kr", "co.in", "co.za",
    "co.nz", "com.au", "net.au", "org.au", "com.br", "com.mx", "com.tr",
```

```
    "com.cn", "com.hk", "com.sg", "com.tw", "com.ar", "co.il", "ne.jp",
}


def sym(ch):
    return AIDX.get(ch, OTHER)


def registered_and_sub(qh):
    """Split a query host into (registered_domain, subdomain_labels_string)."""
    qh = qh.rstrip(".").lower()
    parts = qh.split(".")
    if len(parts) <= 2:
        return qh, ""
    last2 = ".".join(parts[-2:])
    reg_n = 3 if last2 in TWO_LEVEL else 2
    reg = ".".join(parts[-reg_n:])
    sub = ".".join(parts[:-reg_n])
    return reg, sub


def iter_querylog(path, recent=0):
    """Stream (QH, QT) from an AdGuard querylog. recent>0 keeps only the last N."""
    if recent > 0:
        # Bounded tail without loading the file: keep a ring buffer of raw lines.
        from collections import deque
        buf = deque(maxlen=recent)
        with open(path, "r", encoding="utf-8", errors="ignore") as f:
            for line in f:
                if line.strip():
                    buf.append(line)
        lines = buf
    else:
        lines = _line_reader(path)
    for line in lines:
        try:
            rec = json.loads(line)
        except (json.JSONDecodeError, ValueError):
            continue
        qh = rec.get("QH")
        if not qh:
            continue
        qh = qh.strip().lower()
        if qh.endswith("in-addr.arpa") or qh.endswith("ip6.arpa"):
            continue
        yield qh, rec.get("QT", "")


def _line_reader(path):
    with open(path, "r", encoding="utf-8", errors="ignore") as f:
        for line in f:
            if line.strip():
                yield line


#  n-gram model
def label_improbability(label, logp):
    """Mean -log2 P(char|prev) over a label. High = doesn't look like real naming."""
    if not label:
        return 0.0
    prev = BEG
    total = 0.0
    for ch in label:
        s = sym(ch)
        total += logp[prev][s]
        prev = s
    return total / len(label)


def name_score(sub, logp):
    """Worst-label improbability across the subdomain (tunneling hides in one label)."""
    if not sub:
        return 0.0
    return max(label_improbability(lbl, logp) for lbl in sub.split(".") if lbl)


def finalize_logp(counts):
    """Add-1 smoothed transition matrix  -log2 probabilities."""
    logp = [[0.0] * NSYM for _ in range(NSTATES)]
    for a in range(NSTATES):
        row_total = sum(counts[a]) + NSYM  # +NSYM for add-1 smoothing
```

```
        for b in range(NSYM):
            p = (counts[a][b] + 1) / row_total
            logp[a][b] = -math.log2(p)
    return logp


def learn(log_path, out_path):
    counts = [[0] * NSYM for _ in range(NSTATES)]
    reg_counts = defaultdict(int)
    reservoir = []          # sampled subdomain strings for threshold calibration
    RES_MAX = 6000
    seen = 0

    for qh, _qt in iter_querylog(log_path):
        reg, sub = registered_and_sub(qh)
        reg_counts[reg] += 1
        if not sub:
            continue
        for lbl in sub.split("."):
            prev = BEG
            for ch in lbl:
                s = sym(ch)
                counts[prev][s] += 1
                prev = s
        # Reservoir-sample subdomains (unbiased) for a robust score distribution.
        seen += 1
        if len(reservoir) < RES_MAX:
            reservoir.append(sub)
        else:
            j = random.randint(0, seen - 1)
            if j < RES_MAX:
                reservoir[j] = sub

    if seen == 0:
        print("learn: no subdomained queries found  nothing to train on.", file=sys.stderr)
        return 1

    logp = finalize_logp(counts)

    # Robust threshold from the score distribution: median + kMAD (MAD1.4826,
    # so k=6  mean+4 but resistant to the CDN-hash outliers that wreck plain ).
    scores = sorted(name_score(s, logp) for s in reservoir if s)
    median = _median(scores)
    mad = _median(sorted(abs(x - median) for x in scores)) or 1e-9
    threshold = median + 6.0 * mad

    # Data-driven allowlist: the user's frequent registered domains (their normal
    # traffic, incl. their habitual CDNs)  tunneling to a NEW domain isn't here.
    top = sorted(reg_counts.items(), key=lambda kv: kv[1], reverse=True)
    allow_cut = max(20, int(0.002 * seen))   # "frequent" = seen a lot for this net
    allowlist = [r for r, c in top if c >= allow_cut][:400]

    model = {
        "version": 2,
        "counts": counts,              # compact ints; logp recomputed on load
        "threshold": threshold,
        "median": median,
        "mad": mad,
        "allowlist": allowlist,
        "trained_on": seen,
        "unique_registered": len(reg_counts),
    }
    with open(out_path, "w", encoding="utf-8") as f:
        json.dump(model, f, separators=(",", ":"))
    print(f"learn: trained on {seen} subdomained queries across "
          f"{len(reg_counts)} registered domains.")
    print(f"  score median={median:.3f} MAD={mad:.3f}  threshold={threshold:.3f}")
    print(f"  allowlist: {len(allowlist)} frequent domains    {out_path} "
          f"({os.path.getsize(out_path)} bytes)")
    return 0


def _median(xs):
    n = len(xs)
    if n == 0:
        return 0.0
    xs = sorted(xs)
    m = n // 2
    return xs[m] if n % 2 else (xs[m - 1] + xs[m]) / 2.0


```

```
#  scan
SUSPICIOUS_QT = {"TXT", "NULL", "CNAME", "ANY"}   # classic tunneling data channels
SUB_CAP = 64                                       # bound per-domain memory

# High-entropy-but-legit CDN/cloud suffixes: their hostnames LOOK random but are
# structured, not encoded payload. A tiny static safety net beside the data-driven
# frequency allowlist (which only catches domains queried a lot in THIS log).
CDN_SUFFIXES = (
    "nflxvideo.net", "nflxso.net", "googleusercontent.com", "usercontent.goog",
    "cloudfront.net", "akamai.net", "akamaihd.net", "akamaiedge.net", "edgekey.net",
    "edgesuite.net", "fastly.net", "fbcdn.net", "ytimg.com", "ggpht.com", "gvt1.com",
    "gvt2.com", "1e100.net", "azureedge.net", "cloudflare.net", "cloudflarestream.com",
    "digitaloceanspaces.com", "b-cdn.net", "llnwd.net", "wac.edgecastcdn.net", "cdn77.org",
)

_HEX = re.compile(r"^[0-9a-f]{16,}$")
_B32 = re.compile(r"^[a-z2-7]{16,}$")


def looks_encoded(label):
    """True if a label looks like an encoded payload (hex/base32/long base64),
    NOT just a random-ish structured CDN name or a long dictionary word. This is
    the strongest single discriminator between DNS tunneling and legitimate
    high-entropy hostnames."""
    if _HEX.match(label):                              # 16+ hex chars
        return True
    # base32 alphabet [a-z2-7] fully overlaps lowercase letters, so a long word
    # like "visualstudioexptteam" would match  require digits, which real base32
    # payload always carries (~19% of chars) but dictionary words don't.
    if _B32.match(label) and sum(c in "234567" for c in label) >= 2:
        return True
    # long base64-ish: high char diversity AND several digits (excludes words).
    if len(label) >= 20 and len(set(label)) >= 13 and sum(c.isdigit() for c in label) >= 3:
        return True
    return False


def is_allowlisted(reg, allow):
    return reg in allow or any(reg == s or reg.endswith("." + s) for s in CDN_SUFFIXES)


def scan(log_path, model_path, recent, quiet):
    # Fail LOUD, never silently: a corrupt/absent model or an unreadable log must
    # NOT look like a clean result  that would be a security detector reporting
    # "all clear" while blind. Each failure prints to stderr and exits non-zero.
    try:
        with open(model_path, "r", encoding="utf-8") as f:
            model = json.load(f)
        counts = model["counts"]
        threshold = float(model["threshold"])
        if (not isinstance(counts, list) or len(counts) != NSTATES
                or any(not isinstance(r, list) or len(r) != NSYM for r in counts)):
            raise ValueError("bigram matrix has the wrong shape")
    except (OSError, json.JSONDecodeError, KeyError, ValueError, TypeError) as e:
        msg = (f"could not load a usable baseline model at {model_path} "
               f"({type(e).__name__}: {e})  run 'learn' first")
        print(f"scan: FAILED to {msg}.", file=sys.stderr)
        log_fail(msg)
        return 2
    logp = finalize_logp(counts)
    allow = set(model.get("allowlist", []))

    agg = {}   # reg -> [queries, {subdomains cap64}, sum_score, n_scored, susp_qt, maxlen, encoded]
    for qh, qt in iter_querylog(log_path, recent=recent):
        reg, sub = registered_and_sub(qh)
        a = agg.get(reg)
        if a is None:
            a = agg[reg] = [0, set(), 0.0, 0, 0, 0, 0]
        a[0] += 1
        if qt in SUSPICIOUS_QT:
            a[4] += 1
        if sub:
            if len(a[1]) < SUB_CAP:
                a[1].add(sub)
            sc = name_score(sub, logp)
            a[2] += sc
            a[3] += 1
            for l in sub.split("."):
                if len(l) > a[5]:
                    a[5] = len(l)
                if not a[6] and looks_encoded(l):
```

```
                    a[6] = 1

    if not agg:
        msg = (f"the querylog at {log_path} yielded 0 parseable DNS records  "
               f"the detector did NOT run; this is NOT a clean result")
        print(f"scan: FAILED  {msg}.", file=sys.stderr)
        log_fail(msg)
        return 2

    findings = []
    for reg, (q, subs, ssum, sn, susp, maxlen, encoded) in agg.items():
        if sn == 0 or is_allowlisted(reg, allow):
            continue
        mean_score = ssum / sn
        uniq = len(subs)
        improbable = mean_score > threshold
        # Three independent tunneling signatures (any one flags), each tunneling-
        # SHAPED so structured CDN names don't trip it. Encoded-payload volume is
        # primary because encoded labels score like CDN hashes under the bigram
        # model  the improbability threshold alone can't separate them, but their
        # SHAPE (pure hex/base32/base64) and per-domain multiplicity can:
        strong_encoded = encoded and uniq >= 8          # many unique encoded payloads
        txt_abuse = susp >= 5                             # sustained TXT/NULL channel
        improbable_vol = improbable and (uniq >= 8 or maxlen >= 24)
        if strong_encoded or txt_abuse or improbable_vol:
            corro = min(1.0, uniq / 30.0 + susp / 12.0 + encoded * 0.45 + (maxlen >= 40) * 0.2)
            over = min(1.0, max(0.0, mean_score - threshold) / (threshold + 1e-9))
            findings.append((round(0.55 * corro + 0.45 * over, 3), reg, round(mean_score, 2),
                             uniq + (SUB_CAP if uniq >= SUB_CAP else 0), susp, maxlen, q))

    findings.sort(reverse=True)
    if quiet:
        print(f"DNS anomaly scan: {len(findings)} suspicious domain(s) "
              f"(threshold={threshold:.2f}, {len(agg)} domains scanned)")
    else:
        print(f"\n=== DNS anomaly scan  {len(findings)} suspicious domain(s) "
              f"of {len(agg)} scanned (threshold {threshold:.2f}) ===")
        if findings:
            print(f"{'conf':>5}  {'registered domain':32} {'score':>6} {'uniqsub':>7} "
                  f"{'txt/null':>8} {'maxlbl':>6} {'queries':>7}")
            for conf, reg, sc, uniq, susp, maxlen, q in findings[:25]:
                cap = "+" if uniq > SUB_CAP else " "
                print(f"{conf:>5}  {reg:32.32} {sc:>6} {min(uniq, SUB_CAP):>6}{cap} "
                      f"{susp:>8} {maxlen:>6} {q:>7}")
        else:
            print("  clean  no domain combined improbable naming with a volumetric signal.")
    return 1 if findings else 0


#  demo (self-test, no querylog needed)
def demo():
    normal = ["google.com", "api.github.com", "mail.google.com", "cdn.jsdelivr.net",
              "en.wikipedia.org", "outlook.office365.com", "i.redd.it", "apple.com",
              "static.cloudflareinsights.com", "play.googleapis.com"] * 40
    counts = [[0] * NSYM for _ in range(NSTATES)]
    for qh in normal:
        _reg, sub = registered_and_sub(qh)
        for lbl in sub.split("."):
            prev = BEG
            for ch in lbl:
                s = sym(ch)
                counts[prev][s] += 1
                prev = s
    logp = finalize_logp(counts)
    samples = [registered_and_sub(q)[1] for q in normal]
    scores = sorted(name_score(s, logp) for s in samples if s)
    med = _median(scores)
    mad = _median(sorted(abs(x - med) for x in scores)) or 1e-9
    thr = med + 6.0 * mad
    print(f"demo threshold = {thr:.3f}  (median {med:.3f}, MAD {mad:.3f})")
    tests = [("api.github.com", "normal"), ("mail.google.com", "normal"),
             ("A8F93BD72Q9X.attacker.com", "base64 tunnel"),
             ("c732488a09b2e4f6d1.tunnel.evil.org", "hex tunnel"),
             ("kZ9x-Qw82-Lp01-Vn55.dns.io", "high-entropy tunnel")]
    for qh, label in tests:
        _reg, sub = registered_and_sub(qh)
        sc = name_score(sub, logp)
        flag = " anomalous" if sc > thr else " normal"
        print(f"  {flag:14} score={sc:5.2f}  {qh}  ({label})")


```

```
def selftest():
    """Prove the detection logic works end-to-end. Exit 0 = healthy, 3 = BROKEN
    (loud). Lets sweep/CI catch a regression instead of silently trusting a
    detector that no longer detects. Validates BOTH signal paths: the n-gram
    language model (improbability) and the encoding-shape detector."""
    normal = ["google.com", "api.github.com", "mail.google.com", "cdn.jsdelivr.net",
              "en.wikipedia.org", "outlook.office365.com", "i.redd.it", "apple.com"] * 40
    counts = [[0] * NSYM for _ in range(NSTATES)]
    for qh in normal:
        for lbl in registered_and_sub(qh)[1].split("."):
            prev = BEG
            for ch in lbl:
                s = sym(ch)
                counts[prev][s] += 1
                prev = s
    logp = finalize_logp(counts)
    scores = sorted(name_score(registered_and_sub(q)[1], logp) for q in normal
                    if registered_and_sub(q)[1])
    med = _median(scores)
    mad = _median(sorted(abs(x - med) for x in scores)) or 1e-9
    thr = med + 6.0 * mad

    fails = []
    for g in ("mail.google.com", "api.github.com", "cdn.jsdelivr.net"):
        sub = registered_and_sub(g)[1]
        if name_score(sub, logp) > thr or any(looks_encoded(l) for l in sub.split(".")):
            fails.append(f"false-positive on known-good '{g}'")
    # hex/base32 encoded payloads must be caught by the SHAPE detector:
    for enc in ("a8f93bd72c9e4f16.evil.com", "mzxw6ytb2do4zb3q7a.tunnel.io"):
        lbl = registered_and_sub(enc)[1].split(".")[0]
        if not looks_encoded(lbl):
            fails.append(f"encoding detector missed '{lbl}'")
    # an unpronounceable non-encoded label must be caught by the n-gram model:
    imp = registered_and_sub("xqzjvkwbmqpzrfxn.dga.net")[1].split(".")[0]
    if name_score(imp, logp) <= thr:
        fails.append(f"n-gram model missed improbable label '{imp}'")

    if fails:
        print("SELFTEST FAILED  detector is not working:", file=sys.stderr)
        for f in fails:
            print(f"  - {f}", file=sys.stderr)
        log_fail("selftest  detector is not working: " + "; ".join(fails))
        return 3
    print("selftest OK  flags hex/base32 payloads + improbable labels, clears known-good.")
    return 0


def main():
    ap = argparse.ArgumentParser(description="nullexit DNS anomaly detector (report-only)")
    sub = ap.add_subparsers(dest="cmd", required=True)
    pl = sub.add_parser("learn"); pl.add_argument("--log", default=DEFAULT_LOG); pl.add_argument("--out",
     default=DEFAULT_MODEL)
    ps = sub.add_parser("scan")
    ps.add_argument("--log", default=DEFAULT_LOG); ps.add_argument("--model", default=DEFAULT_MODEL)
    ps.add_argument("--recent", type=int, default=0, help="scan only the last N records (0=all)")
    ps.add_argument("--quiet", action="store_true")
    sub.add_parser("demo")
    sub.add_parser("selftest")
    args = ap.parse_args()

    if args.cmd == "learn":
        if not os.path.exists(args.log):
            print(f"learn: querylog not found at {args.log}", file=sys.stderr); return 1
        return learn(args.log, args.out)
    if args.cmd == "scan":
        if not os.path.exists(args.log):
            print(f"scan: querylog not found at {args.log}", file=sys.stderr); return 2
        return scan(args.log, args.model, args.recent, args.quiet)
    if args.cmd == "demo":
        demo(); return 0
    if args.cmd == "selftest":
        return selftest()


if __name__ == "__main__":
    sys.exit(main())
```

39 scripts/fix-docker-bridge-collision.sh

```
#!/bin/bash
# scripts/fix-docker-bridge-collision.sh  nullexit fix for devref 9.13
#
# Symptom: host default route points at Docker's bridge gateway (172.17.0.1)
# instead of the real Wi-Fi gateway. ALL host internet egress fails. Remote
# tailnet clients route through the gateway correctly (their tunnel is
# independent of the host's broken route table), so they look healthy while
# the host itself cannot reach anything.
#
# Root cause: Docker's default bridge claims 172.17.0.0/16 by default. When the
# host's physical Wi-Fi also lands in the same /16  extremely common on
# university networks which hand out 172.16.0.0/12  the kernel
# installs Docker's bridge as the default route, and every packet bound for
# the internet is silently absorbed by docker0 instead of egressing.
#
# This script:
#   1. Detects the actual collision (Docker's 172.17.0.0/16 vs the host's
#      Wi-Fi subnet, queried live via the ACTIVE interface  not hardcoded en0)
#   2. Picks a non-colliding Docker bip using /8-family routing (works for
#      172.16.0.0/12 campus networks; naive prefix-match would pick 172.26/24
#      which still lives INSIDE a /12 of 172.16.0.0/12)
#   3. Backs up ~/.colima/default/colima.yaml with an ISO-8601 timestamp
#   4. Edits the file in place, idempotent (replace existing docker.bip /
#      append new docker.bip without disturbing other keys; skip YAML comments)
#   5. Restarts Colima so the VM rebuilds with the new bridge subnet
#   6. Rebinds Wi-Fi DHCP so the host re-learns its real Wi-Fi gateway
#   7. Verifies the new default route is no longer the Docker bridge
#   8. Re-runs toggle.sh to bring the gateway back up cleanly
#   9. Re-runs scripts/diagnose-host-leak.sh and prints the new verdict
#
# Usage:
#   bash scripts/fix-docker-bridge-collision.sh        # apply the fix
#   bash scripts/fix-docker-bridge-collision.sh --dry  # show what would change, then exit
#   bash scripts/fix-docker-bridge-collision.sh --skip-toggle  # fix but don't restart gateway
#   bash scripts/fix-docker-bridge-collision.sh --help # usage

set -uo pipefail
# NOTE: errexit (`set -e`) is intentionally NOT enabled so intermediate
# `command  || warn ` patterns can continue past non-fatal failures.
# Every critical failure (colima restart, sed, cp backup, toggle.sh)
# uses an explicit `|| die` so a half-applied state never escapes.

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
LOG_FILE="$PROJECT_ROOT/output.log"
source "$SCRIPT_DIR/common.sh"
TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
COLIMA_YAML="$HOME/.colima/default/colima.yaml"
PLATFORM="$(uname -s)"

#  Argument parsing
DRY_RUN=false
SKIP_TOGGLE=false
case "${1:-}" in
  --dry|-n)       DRY_RUN=true ;;
  --skip-toggle)  SKIP_TOGGLE=true ;;
  --help|-h)
    sed -n '28,37p' "$0"
    exit 0
    ;;
  "")
    : # default: apply
    ;;
  *)
    printf 'Unknown argument: %s\n' "$1" >&2
    exit 2
    ;;
esac



# Helper: print "[dry-run] $cmd" or actually run $cmd. Wraps the redirect
# INSIDE the function so outer `>> $FILE` statements never accidentally
# append dry-run noise to the real file. (Earlier draft didn't and the
# dry-run polluted $COLIMA_YAML.)
run() {
  if [ "$DRY_RUN" = "true" ]; then
```

```
    printf '    [dry-run] %s\n' "$*"
    return 0
  fi
  printf '    $ %s\n' "$*"
  "$@"
}

# Platform-aware sed in-place. macOS BSD sed needs an empty argument;
# GNU sed (Linux) does not. Earlier draft used `-i ''` everywhere and
# would have failed on Linux silently.
sed_inplace() {
  local expr="$1"
  local file="$2"
  case "$PLATFORM" in
    Darwin) run sed -i '' "$expr" "$file" ;;
    *)      run sed -i    "$expr" "$file" ;;
  esac
}

#  0. Resolve the ACTIVE physical interface once (used everywhere)
# Replaces the earlier `en0` hardcoding. Same algorithm recover.sh uses:
# `route get default`  if utun/tun, walk en0..en5 first with `status: active`
# + `inet`  fall back to en0. This means the script also works on hosts
# whose primary NIC is en1/en2 (Thunderbolt Ethernet, USB-LAN, etc.).
resolve_active_iface() {
  local iface
  iface=$(route get default 2>> "$LOG_FILE" | awk '/interface:/{print $2; exit}')
  if [ -z "$iface" ] || [ "${iface#utun}" != "$iface" ] || [ "${iface#tun}" != "$iface" ]; then
    for i in en0 en1 en2 en3 en4 en5; do
      if ifconfig "$i" 2>> "$LOG_FILE" | grep -q "status: active" \
         && ifconfig "$i" 2>> "$LOG_FILE" | grep -q "inet "; then
        iface="$i"; break
      fi
    done
  fi
  [ -z "$iface" ] && iface="en0"
  printf '%s\n' "$iface"
}
ACTIVE_IFACE=$(resolve_active_iface)

#  1. Detect collision
step "1/9  Detecting Docker-bridge vs Wi-Fi subnet collision"

# Host subnet on the ACTIVE interface (not hardcoded en0).
HOST_SUBNET=$(ifconfig "$ACTIVE_IFACE" 2>> "$LOG_FILE" \
              | awk '/inet /{print $2; exit}')
HOST_GATEWAY=$(route get default 2>> "$LOG_FILE" \
               | awk '/gateway:/{print $2; exit}')
DEFAULT_DEV=$(route get default 2>> "$LOG_FILE" \
              | awk '/interface:/{print $2; exit}')

printf '  active iface:  %s\n' "$ACTIVE_IFACE"
printf '  iface inet:    %s\n' "${HOST_SUBNET:-none}"
printf '  default via:   %s\n' "${HOST_GATEWAY:-none}"
printf '  default dev:   %s\n' "${DEFAULT_DEV:-none}"

# Collision fires when the host's default-route gateway is 172.17.0.1 (Docker's
# bridge) OR when the host subnet is in 172.17.0.0/16. The full-bridge-collision
# case (gateway is literally the docker bridge IP) is the smoking gun.
COLLISION=false
if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#172.17.}" != "$HOST_GATEWAY" ]; then
  COLLISION=true
  fail "default gateway $HOST_GATEWAY collides with Docker's default bridge (172.17.0.0/16)"
elif [ -n "$HOST_SUBNET" ] && [ "${HOST_SUBNET#172.17.}" != "$HOST_SUBNET" ]; then
  COLLISION=true
  fail "host subnet $HOST_SUBNET overlaps Docker's 172.17.0.0/16 bridge"
else
  ok "no collision detected between host Wi-Fi and Docker's default bridge"
fi

if [ "$COLLISION" = "false" ]; then
  echo ""
  echo "Nothing to fix. If you still see egress issues, run:"
  echo "    bash scripts/diagnose-host-leak.sh"
  exit 0
fi

#  2. Pick non-colliding bip via /8-family routing
step "2/9  Picking a non-colliding Docker bip for ~/.colima/default/colima.yaml"

# Earlier draft tried naive prefix-match: "if candidate's 3-octet prefix is
```

```
# not a literal prefix of host gateway, pick it". That breaks on /12 campus
# networks (172.16.0.0/12 covers 172.16-172.31): a candidate like
# 172.26.0.1/24 LIVES INSIDE the host's /12 even though the prefix test
# passes. Fix: classify the host's address family FIRST, then choose ALL
# candidates from a DIFFERENT RFC1918 family. /8 boundaries are coarse
# enough that no /24 picked this way can ever collide.
case "$HOST_GATEWAY" in
  172.16.*|172.17.*|172.18.*|172.19.*|172.20.*|172.21.*|172.22.*|172.23.*|\
  172.24.*|172.25.*|172.26.*|172.27.*|172.28.*|172.29.*|172.30.*|172.31.*)
    HOST_FAMILY="172"
    # Host is in 172.x  jump past 172.31 entirely. Pick from 10.x / 192.168.x.
    CANDIDATES="10.200.0.1/24 10.201.0.1/24 192.168.99.1/24"
    ;;
  10.*)
    HOST_FAMILY="10"
    # Host is in 10.x  pick from 172.x (well outside any commonly-deployed
    # 10.0.0.0/8 slice) / 192.168.x.
    CANDIDATES="172.26.0.1/24 172.28.0.1/24 192.168.99.1/24"
    ;;
  192.168.*)
    HOST_FAMILY="192"
    CANDIDATES="172.26.0.1/24 10.200.0.1/24 10.201.0.1/24"
    ;;
  *)
    HOST_FAMILY="unknown"
    warn "host gateway $HOST_GATEWAY is not in any RFC1918 range (/8 family)"
    warn "falling back to abstract candidate set; verify collision-free manually"
    CANDIDATES="172.26.0.1/24 10.200.0.1/24 192.168.99.1/24"
    ;;
esac

printf '  host address family: %s\n' "$HOST_FAMILY"
printf '  candidate pool:      %s\n' "$CANDIDATES"

# Sanity-check the chosen pool has at least one entry that doesn't literally
# share a /24 with the host gateway. Naive check (just /24 prefix) is fine
# here because we already picked across families.
CHOSEN=""
for cand in $CANDIDATES; do
  PREFIX3=$(printf '%s' "$cand" | cut -d. -f1-3)
  if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#"$PREFIX3".}" != "$HOST_GATEWAY" ]; then
    warn "skip $cand (literal /24 collision with host gateway $HOST_GATEWAY)"
    continue
  fi
  CHOSEN="$cand"
  ok "chose bip=$CHOSEN (cross-family, no /24 literal collision)"
  break
done

if [ -z "$CHOSEN" ]; then
  CHOSEN="172.26.0.1/24"
  warn "could not find a clean candidate; defaulting to $CHOSEN"
  warn "verify with 'route get default' after restart"
fi

#  3. Backup colima.yaml
step "3/9  Backing up $COLIMA_YAML"
mkdir -p "$(dirname "$COLIMA_YAML")"
BACKUP="$HOME/.colima/default/colima.yaml.bak.$TIMESTAMP"
if [ -f "$COLIMA_YAML" ]; then
  # `cp` is safe in dry-run because run() suppresses the inner command
  # entirely  no chance of accidentally writing a real backup.
  run cp -p "$COLIMA_YAML" "$BACKUP" || die "backup failed"
  if [ "$DRY_RUN" = "false" ]; then ok "backup: $BACKUP"; fi
else
  warn "no existing $COLIMA_YAML (will create fresh)"
fi

#  4. Edit colima.yaml in place (idempotent)
step "4/9  Setting docker.bip=$CHOSEN in colima.yaml (in place)"

# Resolve the existing docker.bip (if any) without picking up commented-out
# YAML keys. awk-based extraction; the negated-comment check (`! /^[[:space:]]*#`)
# ensures `# bip: 1.2.3.4/24` is treated as a comment, not the active value.
EXISTING_BIP=""
if [ -f "$COLIMA_YAML" ]; then
  EXISTING_BIP=$(awk '
    /^[[:space:]]*#/ {next}
    /^[[:space:]]*bip:[[:space:]]*/ {print $2; exit}
  ' "$COLIMA_YAML" 2>> "$LOG_FILE" || true)
fi
```

```
240
241 if [ -n "$EXISTING_BIP" ] && [ "$EXISTING_BIP" = "$CHOSEN" ]; then
242   ok "colima.yaml already has docker.bip=$CHOSEN  no edit needed"
243 elif [ -n "$EXISTING_BIP" ]; then
244   printf '  (replacing existing docker.bip=%s with %s)\n' "$EXISTING_BIP" "$CHOSEN"
245   # Use sed_inplace so Linux / macOS get the right syntax. Skip YAML comment
246   # lines so existing `# bip:` comments aren't rewired.
247   sed_inplace "s|^\\([[:space:]]*bip:[[:space:]]*\\).*\$|\\1$CHOSEN|" "$COLIMA_YAML" \
248     || die "sed replacement failed (colima.yaml malformed?)"
249   ok "colima.yaml updated"
250 else
251   # No existing `bip:` key. Branch:
252   #   - file has `docker:` block  append `  bip:` under it
253   #   - file has no `docker:` block  append new docker: section
254   #   - file does not exist  create fresh
255   # Each branch is dry-run-aware so no branch can accidentally leak noise
256   # into $COLIMA_YAML during a preview.
257   HAS_DOCKER_BLOCK=false
258   if [ -f "$COLIMA_YAML" ] && grep -q '^docker:' "$COLIMA_YAML"; then
259     HAS_DOCKER_BLOCK=true
260   fi
261
262   if [ "$HAS_DOCKER_BLOCK" = "true" ]; then
263     if [ "$DRY_RUN" = "true" ]; then
264       printf '    [dry-run] would append to %s:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
265     else
266       printf '\n  bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
267         || die "append failed"
268       ok "appended bip under existing docker: block"
269     fi
270   elif [ -f "$COLIMA_YAML" ]; then
271     if [ "$DRY_RUN" = "true" ]; then
272       printf '    [dry-run] would append to %s:\\ndocker:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
273     else
274       printf '\ndocker:\n  bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
275         || die "append failed"
276       ok "appended new docker: block at end of file"
277     fi
278   else
279     if [ "$DRY_RUN" = "true" ]; then
280       printf '    [dry-run] would create %s with:\\ndocker:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
281     else
282       printf 'docker:\n  bip: %s\n' "$CHOSEN" > "$COLIMA_YAML" \
283         || die "create failed"
284       ok "created fresh $COLIMA_YAML with docker.bip=$CHOSEN"
285     fi
286   fi
287 fi
288
289 # Always show the final colima.yaml docker section for visibility
290 printf '\n%b  colima.yaml `docker:` section now reads:%b\n' "$BOLD" "$NC"
291 awk '
292   /^docker:/{p=1}
293   p && /^[^ ]/ && !/^docker:/{exit}
294   p
295 ' "$COLIMA_YAML" 2>/dev/null | sed 's/^/    /' || true
296
297 if [ "$DRY_RUN" = "true" ]; then
298   printf '\n  %b(dry-run: exiting before colima restart / dhcp rebind / toggle.sh)%b\n' "$YELLOW" "$NC"
299   exit 0
300 fi
301
302 #  5. Restart Colima
303 step "5/9  Restarting Colima VM with new bip"
304 if ! command -v colima >> "$LOG_FILE" 2>&1; then
305   die "colima not on PATH  install with 'brew install colima' first"
306 fi
307 # No timeout wrapper (would mask actual VM startup time). If colima hangs,
308 # the user kills with Ctrl+C and can manually `colima restart` from a fresh
309 # terminal  the rest of this script is idempotent so re-running works.
310 colima restart >> "$LOG_FILE" 2>&1 || die "colima restart failed; check $LOG_FILE"
311 ok "colima restarted"
312
313 #  6. Rebind DHCP on the ACTIVE interface (not hardcoded en0)
314 step "6/9  Rebinding DHCP on $ACTIVE_IFACE so host re-learns its real gateway"
315
316 # Map ifname  macOS network service name (Wi-Fi, USB 10/100 LAN, etc.).
317 # Same logic as the diagnostic + toggle.sh.
318 ACTIVE_SVC=""
319 if [ "$PLATFORM" = "Darwin" ]; then
320   ACTIVE_SVC=$(get_active_service)
```

```
fi

case "$PLATFORM" in
  Darwin)
    sudo -n networksetup -setdhcp "$ACTIVE_SVC" renew >> "$LOG_FILE" 2>&1 \
      && ok "DHCP rebind issued on $ACTIVE_SVC" \
      || warn "setdhcp renew failed; try 'sudo ipconfig set $ACTIVE_IFACE DHCP' manually"
    ;;
  Linux)
    sudo -n dhclient -r "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 || true
    sudo -n dhclient "$ACTIVE_IFACE"    >> "$LOG_FILE" 2>&1 \
      && ok "dhclient rebind on $ACTIVE_IFACE" \
      || warn "dhclient rebind failed"
    ;;
  *)
    warn "unsupported platform $PLATFORM; skip DHCP rebind"
    ;;
esac

#  7. Verify new default route
step "7/9  Verifying the new default route is NOT the Docker bridge"
sleep 2  # give DHCP + kernel a beat
NEW_DEFAULT=$(netstat -rn 2>> "$LOG_FILE" | awk '/^default/ {print $0; exit}')
printf '  %s\n' "${NEW_DEFAULT:-no default route detected}"
NEW_GATEWAY=$(printf '%s' "$NEW_DEFAULT" | awk '{print $2; exit}')

if [ -n "$NEW_GATEWAY" ] && [ "${NEW_GATEWAY#172.17.}" != "$NEW_GATEWAY" ]; then
  fail "default route STILL points at 172.17.x.x  DHCP rebind may need manual touch"
  warn "try: sudo ipconfig set $ACTIVE_IFACE DHCP  (or reconnect Wi-Fi in the menu bar)"
elif [ -n "$NEW_GATEWAY" ]; then
  ok "default route is now via $NEW_GATEWAY (no longer Docker's bridge)"
else
  warn "could not determine new gateway; verify with 'route get default'"
fi

#  8. Re-run toggle.sh to bring the gateway back up
if [ "$SKIP_TOGGLE" = "true" ]; then
  step "8/9  Skipping toggle.sh (--skip-toggle)"
  warn "gateway is currently down  run ./toggle.sh manually when ready"
else
  step "8/9  Re-running toggle.sh to bring the gateway back up"
  if bash "$PROJECT_ROOT/toggle.sh" 2>&1 | tee -a "$LOG_FILE"; then
    ok "toggle.sh completed"
  else
    warn "toggle.sh returned non-zero  see $LOG_FILE for details"
  fi
fi

#  9. Re-run the diagnostic + show verdict
step "9/9  Re-running host-leak diagnostic to confirm fix"
if bash "$SCRIPT_DIR/diagnose-host-leak.sh" 2>&1 | tee -a "$LOG_FILE"; then
  LATEST_REPORT=$(ls -t "$PROJECT_ROOT"/host-leak-diagnostic-*.txt 2>/dev/null | head -1)
  if [ -n "$LATEST_REPORT" ]; then
    VERDICT=$(awk '/^ Scenario:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
    WARP=$(awk -F': *' '/^ WARP egress:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
    printf '\n  Most recent verdict: Scenario=%s  WARP=%s\n' \
           "${VERDICT:-?}" "${WARP:-?}"
    if [ "$VERDICT" = "OK" ]; then
      ok "fix confirmed by diagnostic  host traffic is going through WARP"
    fi
  fi
else
  warn "diagnostic exited non-zero; verify manually"
fi

printf '\n%b%b\n' "$BOLD" "$NC"
printf '%bDone.%b Full transcript at %s\n' "$BOLD" "$NC" "$LOG_FILE"
printf '%b%b\n' "$BOLD" "$NC"
exit 0
```

40 scripts/fix-ssh-delay.sh

```
#!/bin/bash
# scripts/fix-ssh-delay.sh  Fix ~20s SSH connection delay caused by UseDNS
#
# macOS OpenSSH defaults UseDNS to "yes". When Tailscale SSH connects from
```

```
# a peer (phone/laptop), sshd performs a reverse-DNS (PTR) lookup on the
# client's 100.x.x.x Tailscale IP. Since DNS is hijacked to the nullexit
# gateway  AdGuard  Cloudflare WARP, and 100.64.0.0/10 is CGNAT private
# space, the PTR query times out (~15-20s) before SSH proceeds.
#
# This script uncomments "UseDNS no" in /etc/ssh/sshd_config and reloads
# sshd. Existing SSH sessions are NOT dropped.
#
# Usage:
#   sudo bash scripts/fix-ssh-delay.sh

set -euo pipefail

SSHD_CONFIG="/etc/ssh/sshd_config"

#  Check we're root
if [ "$(id -u)" -ne 0 ]; then
    echo "ERROR: This script must be run with sudo (it edits $SSHD_CONFIG)"
    echo "Usage: sudo bash scripts/fix-ssh-delay.sh"
    exit 1
fi

#  Apply config (idempotent)
if grep -qE '^UseDNS no' "$SSHD_CONFIG" 2>/dev/null; then
    echo " UseDNS is already set to 'no' in $SSHD_CONFIG"
elif grep -qE '^#UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
    echo " Uncommenting 'UseDNS no' in $SSHD_CONFIG"
    sed -i '' 's/^#UseDNS no/UseDNS no/' "$SSHD_CONFIG"
elif grep -qEi '^UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
    echo " Changing existing UseDNS line to 'UseDNS no'"
    sed -i '' 's/^UseDNS.*/UseDNS no/' "$SSHD_CONFIG"
else
    echo " Appending 'UseDNS no' to $SSHD_CONFIG"
    echo "" >> "$SSHD_CONFIG"
    echo "UseDNS no" >> "$SSHD_CONFIG"
fi

#  Verify
if ! grep -qE '^UseDNS no' "$SSHD_CONFIG"; then
    echo " Failed to apply UseDNS no  check $SSHD_CONFIG manually"
    exit 1
fi

#  Reload sshd (does NOT drop existing connections)
echo " Reloading sshd..."

SSHD_PID=$(pgrep -f /usr/sbin/sshd 2>/dev/null | head -1 || true)
if [ -n "$SSHD_PID" ]; then
    kill -HUP "$SSHD_PID"
    echo " Sent SIGHUP to sshd (PID $SSHD_PID)  config reloaded"
else
    echo "  No running sshd found. macOS starts sshd on-demand per connection."
    echo "  The new 'UseDNS no' config will take effect on the next SSH connection."
fi

echo ""
echo " Done. SSH connections should now connect instantly."
echo "  Existing sessions were not interrupted."
```

41 scripts/generate_tex.py

```
#!/usr/bin/env python3
import os

PROJECT_ROOT = os.path.abspath(os.path.join(os.path.dirname(__file__), '..'))

IGNORE_DIRS = {'.git', '.github', 'adguard', '__pycache__', 'Recover Gateway.app', 'Toggle Gateway.app', 'Sweep
      Gateway.app','nullexit-knowledge-graph'}
IGNORE_FILES = {'.env', 'nullexit_unified.tex', 'nullexit_unified.pdf', 'output.log', 'blocked.log', '.DS_Store
     ', '.lan_p2p_detected', '.signatures', '.gateway_ip', 'TUNNEL_FAILED_CLOSED.marker','.host_ips','refactor.
     md','sweep.md','.build_hash','.rules_hash','knowledge-graph.html','.dns_baseline.json','.last_wifi_ssid','.
     last_wifi_ssid.tmp'}
IGNORE_EXTS = {'.pdf', '.tex', '.png', '.jpg', '.jpeg', '.zip', '.tar', '.gz', '.log', '.aux', '.fls', '.
     fdb_latexmk', '.out', '.toc', '.xdv', '.synctex(busy)'}
# Transient runtime reports (timestamped names) carry live egress/peer IPs  never
# embed them in the published quine. Matched by filename prefix.
```

```
IGNORE_PREFIXES = ('sweep-', 'host-leak-diagnostic-')

# Exclude from code block inclusion, but keep in tree
SKIP_CONTENT_FILES = {'LICENSE'}

PRIORITY_FILES = [
    'README.md',
    'agent.md',
    'devref.md',
    'diagrams.md',
    'toggle.sh',
    'recover.sh',
    'setup.sh',
    'scripts/setup-linux.sh',
    'docker-compose.yml',
    'scripts/common.sh',
    'scripts/watcher.sh',
    'scripts/crypto.sh'
]

def get_language(filename):
    ext = os.path.splitext(filename)[1].lower()
    if ext == '.sh' or ext == '.applescript':
        return 'bash'
    elif ext == '.py':
        return 'Python'
    elif ext == '.go':
        return 'Go'
    elif ext in {'.yml', '.yaml'}:
        return 'bash'
    elif ext == '.plist':
        return 'XML'
    elif ext == '.md':
        return ''
    return ''

def escape_tex(text):
    return text.replace('_', '\\_')

def generate_tree(dir_path, prefix=""):
    tree_str = ""
    entries = sorted(os.listdir(dir_path))
    entries = [e for e in entries if e not in IGNORE_DIRS and e not in IGNORE_FILES and os.path.splitext(e)[1].
     lower() not in IGNORE_EXTS]

    for i, entry in enumerate(entries):
        path = os.path.join(dir_path, entry)
        is_last = (i == len(entries) - 1)
        connector = "`-- " if is_last else "|-- "
        tree_str += f"{prefix}{connector}{entry}\n"

        if os.path.isdir(path):
            extension = "    " if is_last else "|   "
            tree_str += generate_tree(path, prefix + extension)

    return tree_str

def main():
    tex_path = os.path.join(PROJECT_ROOT, 'nullexit_unified.tex')

    preamble = r"""\documentclass{article}
\usepackage[utf8]{inputenc}
\usepackage[margin=1in]{geometry}
\usepackage{listings}
\usepackage{xcolor}
\usepackage{hyperref}
\usepackage{tocloft}

\definecolor{codegreen}{rgb}{0,0.6,0}
\definecolor{codegray}{rgb}{0.5,0.5,0.5}
\definecolor{codepurple}{rgb}{0.58,0,0.82}
\definecolor{backcolour}{rgb}{0.97,0.97,0.96}

\lstdefinestyle{mystyle}{
    backgroundcolor=\color{backcolour},
    commentstyle=\color{codegreen},
    keywordstyle=\color{blue}\bfseries,
    numberstyle=\tiny\color{codegray},
    stringstyle=\color{codepurple},
    basicstyle=\ttfamily\scriptsize,
    breakatwhitespace=false,
```

```
    breaklines=true,
    captionpos=b,
    keepspaces=true,
    numbers=left,
    numbersep=5pt,
    showspaces=false,
    showstringspaces=false,
    showtabs=false,
    tabsize=2,
    frame=single,
    rulecolor=\color{codegray},
    extendedchars=true,
    literate={}{-}1 {}{=}1 {}{{OK}}2 {}{-}1 {}{|}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{v}1
}
\lstset{style=mystyle}

\title{Nullexit: Unified Project Document}
\author{Omar Alaaeldein}
\date{\today}

\begin{document}
\maketitle

\tableofcontents
\newpage

\section{Project Directory Tree}
\begin{verbatim}
nullexit/
"""

    tree = generate_tree(PROJECT_ROOT)
    midamble = r"""\end{verbatim}
\newpage
"""

    all_relpaths = []
    for root, dirs, files in os.walk(PROJECT_ROOT):
        dirs[:] = sorted([d for d in dirs if d not in IGNORE_DIRS])
        for file in sorted(files):
            if file in IGNORE_FILES or file in SKIP_CONTENT_FILES or os.path.splitext(file)[1].lower() in IGNORE_EXTS:
                continue
            if file.startswith(IGNORE_PREFIXES):
                continue

            filepath = os.path.join(root, file)
            relpath = os.path.relpath(filepath, PROJECT_ROOT)
            all_relpaths.append(relpath)

    def sort_key(path):
        if path in PRIORITY_FILES:
            return (0, PRIORITY_FILES.index(path))
        return (1, path)

    all_relpaths.sort(key=sort_key)

    with open(tex_path, 'w', encoding='utf-8') as f:
        f.write(preamble)
        f.write(tree)
        f.write(midamble)

        for relpath in all_relpaths:
            lang = get_language(relpath)
            lang_attr = f"[language={lang}]" if lang else ""

            f.write(f"\\section{{{escape_tex(relpath)}}}\n")
            f.write(f"\\begin{{lstlisting}}{lang_attr}\n")
            try:
                with open(os.path.join(PROJECT_ROOT, relpath), 'r', encoding='utf-8', errors='ignore') as src:
                    content = src.read()
                    import re
                    # Strip all non-ASCII characters (emojis, box drawings, em dashes, etc)
                    # AND ASCII control characters except \t\n (binary payloads, ANSI escapes,
                    # form feeds  all invisible, all fatal to LaTeX compilation)
                    content = re.sub(r'[^\x09\x0A\x20-\x7E]+', '', content)
                    f.write(content)
                    if not content.endswith('\n'):
                        f.write('\n')
            except Exception as e:
```

```
                f.write(f"Error reading file: {e}\n")
            f.write("\\end{lstlisting}\n\n")

        f.write("\\end{document}\n")

    print(f"Generated {tex_path}")

if __name__ == '__main__':
    main()
```

42 scripts/logger/go.mod

```
module nullexit/logger

go 1.22
```

43 scripts/logger/main.go

```
package main

import (
  "bufio"
  "encoding/json"
  "fmt"
  "io"
  "log"
  "os"
  "regexp"
  "strings"
  "time"
)

const (
  queryLogPath  = "/adguard_work/data/querylog.json"
  procKmsgPath  = "/proc/kmsg"
  outputLogPath = "/app/blocked.log"
)

func logMessage(msg string) {
  // Always log in UTC with an explicit marker. The routing-fix image is
  // Alpine (no tzdata), so time.Now() already fell back to UTC  but pin it
  // so the format can't silently flip to local time if tzdata is ever added.
  // AdGuard's querylog.json is now UTC too (TZ=UTC in docker-compose.yml).
  timestamp := time.Now().UTC().Format("2006-01-02 15:04:05") + " UTC"
  formattedMsg := fmt.Sprintf("%s %s\n", timestamp, msg)
  fmt.Print(formattedMsg)

  f, err := os.OpenFile(outputLogPath, os.O_APPEND|os.O_CREATE|os.O_WRONLY, 0644)
  if err != nil {
    log.Printf("Error writing to log file: %v\n", err)
    return
  }
  defer f.Close()
  f.WriteString(formattedMsg)
}

func tailFile(filepath string, out chan<- string) {
  for {
    if _, err := os.Stat(filepath); os.IsNotExist(err) {
      time.Sleep(1 * time.Second)
      continue
    }
    break
  }

  f, err := os.Open(filepath)
  if err != nil {
    log.Printf("Error opening file %s: %v\n", filepath, err)
    return
  }
  defer func() {
    if f != nil {
```

```
      f.Close()
    }
  }()

  // Seek to end
  f.Seek(0, io.SeekEnd)
  reader := bufio.NewReader(f)

  for {
    line, err := reader.ReadString('\n')
    if err != nil {
      if err == io.EOF {
        stat, err := os.Stat(filepath)
        if err == nil {
          currentOffset, _ := f.Seek(0, io.SeekCurrent)
          if stat.Size() < currentOffset {
            // File truncated or rotated
            f.Close()
            for {
              if _, err := os.Stat(filepath); os.IsNotExist(err) {
                time.Sleep(1 * time.Second)
                continue
              }
              break
            }
            f, _ = os.Open(filepath)
            reader = bufio.NewReader(f)
            continue
          }
        }
        time.Sleep(500 * time.Millisecond)
        continue
      }
      log.Printf("Error reading file %s: %v\n", filepath, err)
      time.Sleep(1 * time.Second)
      continue
    }
    out <- line
  }
}

type adguardLogEntry struct {
  IP     string `json:"IP"`
  QH     string `json:"QH"`
  QT     string `json:"QT"`
  Result struct {
    IsFiltered bool `json:"IsFiltered"`
    Reason     int  `json:"Reason"`
    Rules      []struct {
      Text string `json:"Text"`
    } `json:"Rules"`
  } `json:"Result"`
}

func monitorDNS() {
  logMessage("[System] Starting DNS block logger...")
  lines := make(chan string)
  go tailFile(queryLogPath, lines)

  reasonMap := map[int]string{
    2: "CustomRule",
    3: "BlockList",
    4: "SafeBrowsing",
    5: "ParentalControl",
    6: "SafeSearch",
    7: "BlockedService",
  }

  for line := range lines {
    var data adguardLogEntry
    if err := json.Unmarshal([]byte(line), &data); err != nil {
      continue
    }

    res := data.Result
    if res.IsFiltered || (res.Reason != 0 && res.Reason != 1) {
      qh := data.QH
      if qh == "" {
        qh = "unknown"
      }
      qt := data.QT
```

```
      if qt == "" {
        qt = "unknown"
      }
      clientIP := data.IP
      if clientIP == "" {
        clientIP = "unknown"
      }

      ruleText := "unknown rule"
      if len(res.Rules) > 0 && res.Rules[0].Text != "" {
        ruleText = res.Rules[0].Text
      }

      reasonStr, ok := reasonMap[res.Reason]
      if !ok {
        reasonStr = fmt.Sprintf("Reason-%d", res.Reason)
      }

      logMessage(fmt.Sprintf("[DNS] Blocked %s (Type: %s) for client %s | Reason: %s | Rule: %s", qh, qt,
     clientIP, reasonStr, ruleText))
    }
  }
}

func monitorIPs() {
  logMessage("[System] Starting IP block logger...")

  prefixRe := regexp.MustCompile(`(IP_BLOCK_[A-Z_]+):`)
  srcRe := regexp.MustCompile(`SRC=([0-9a-fA-F\.:]+)`)
  dstRe := regexp.MustCompile(`DST=([0-9a-fA-F\.:]+)`)
  protoRe := regexp.MustCompile(`PROTO=([A-Z0-9]+)`)
  sptRe := regexp.MustCompile(`SPT=(\d+)`)
  dptRe := regexp.MustCompile(`DPT=(\d+)`)
  inRe := regexp.MustCompile(`IN=([a-zA-Z0-9\.\-_]+)`)
  outRe := regexp.MustCompile(`OUT=([a-zA-Z0-9\.\-_]+)`)

  f, err := os.Open(procKmsgPath)
  if err != nil {
    if os.IsPermission(err) {
      logMessage("[System] ERROR: Insufficient permissions to read /proc/kmsg. Enable CAP_SYSLOG.")
    } else {
      logMessage(fmt.Sprintf("[System] ERROR in IP logger: %v", err))
    }
    return
  }
  defer f.Close()

  reader := bufio.NewReader(f)
  for {
    line, err := reader.ReadString('\n')
    if err != nil {
      if err == io.EOF {
        time.Sleep(100 * time.Millisecond)
        continue
      }
      logMessage(fmt.Sprintf("[System] ERROR reading kmsg: %v", err))
      time.Sleep(1 * time.Second)
      continue
    }

    if strings.Contains(line, "IP_BLOCK") {
      match := prefixRe.FindStringSubmatch(line)
      if len(match) > 1 {
        prefix := match[1]

        src := "unknown"
        if m := srcRe.FindStringSubmatch(line); len(m) > 1 {
          src = m[1]
        }

        dst := "unknown"
        if m := dstRe.FindStringSubmatch(line); len(m) > 1 {
          dst = m[1]
        }

        proto := "unknown"
        if m := protoRe.FindStringSubmatch(line); len(m) > 1 {
          proto = m[1]
        }

        spt := ""
```

```
        if m := sptRe.FindStringSubmatch(line); len(m) > 1 {
          spt = m[1]
        }

        dpt := ""
        if m := dptRe.FindStringSubmatch(line); len(m) > 1 {
          dpt = m[1]
        }

        inIf := ""
        if m := inRe.FindStringSubmatch(line); len(m) > 1 {
          inIf = m[1]
        }

        outIf := ""
        if m := outRe.FindStringSubmatch(line); len(m) > 1 {
          outIf = m[1]
        }

        direction := "inbound"
        if strings.HasSuffix(prefix, "DST") {
          direction = "outbound"
        }

        listType := "MALICIOUS"
        if !strings.Contains(prefix, "MALICIOUS") {
          parts := strings.Split(prefix, "_")
          if len(parts) >= 3 {
            listType = "GEO_" + parts[2]
          }
        }

        portStr := ""
        if dpt != "" {
          portStr = ":" + dpt
        }

        srcPortStr := ""
        if spt != "" {
          srcPortStr = ":" + spt
        }

        ifStr := ""
        if inIf != "" && outIf != "" {
          ifStr = fmt.Sprintf("IF: %s->%s", inIf, outIf)
        } else if inIf != "" || outIf != "" {
          ifStr = fmt.Sprintf("IF: %s%s", inIf, outIf)
        }

        logMessage(fmt.Sprintf("[IP] Blocked %s to %s%s (Proto: %s) from %s%s (%s) | List: %s", direction, dst,
      portStr, proto, src, srcPortStr, ifStr, listType))
      }
    }
  }
}

func main() {
  logMessage("[System] Nullexit Block Logger starting...")

  go monitorDNS()
  monitorIPs()
}
```

44 scripts/noise.sh

```
#!/bin/bash
# scripts/noise.sh  nullexit verifiable-random noise + cover-traffic padding
#
# OPT-IN and ACTIVE. Unlike sweep.sh (read-only), this module PUTS TRAFFIC ON
# THE WIRE. It stays completely inert unless NOISE_ENABLED=true in .env, so the
# default posture is unchanged and the sweep's read-only guarantee is intact.
#
# Two jobs, one primitive (CSPRNG bytes from the OS, /dev/urandom via os.urandom):
#
#   1. verify  prove a freshly generated buffer is statistically random using
#      three standard tests (Shannon entropy, NIST monobit, chi-square over the
```

```
#      256 byte values). This is the "we can verify it's truly random
#      mathematically" requirement  it is harmless and always allowed.
#
#   2. pad  a cover-traffic daemon. It emits verified-random UDP datagrams
#      toward the exit at a configured rate + jitter so a passive on-path
#      observer (residence / campus / hostile network) sees a continuously
#      varying encrypted flow instead of your real traffic envelope. This is
#      the "dummy padding" defence the Threat Model flags as the countermeasure
#      to metadata / traffic-flow correlation. It ONLY runs with NOISE_ENABLED=true.
#
# HONEST LIMITS: this is constant/random one-way padding, not a full
# traffic-analysis-resistant shaper. It blurs volume + timing on the uplink; it
# does not mimic a specific protocol, pad bidirectionally, or defeat a global
# active adversary. It costs bandwidth and (on battery devices) power  hence
# opt-in, modest default rate. The DNS/IP anomaly detector is deliberately NOT
# here (left for later).
#
# Usage:
#   bash scripts/noise.sh verify [BYTES]   # randomness self-test (default 1 MiB)
#   bash scripts/noise.sh start            # launch the padding daemon (background)
#   bash scripts/noise.sh pad              # run the padding loop in the foreground
#   bash scripts/noise.sh status           # is the daemon running? target + rate
#   bash scripts/noise.sh stop             # stop the daemon
#   bash scripts/noise.sh --help           # this message
#
# Persistence: install launchd/com.nullexit.noise.plist (see README) so padding
# auto-starts at boot/login. It is governed by the SAME NOISE_ENABLED switch  no
# extra knob: the launchd job is a clean no-op whenever NOISE_ENABLED=false.
#
# .env knobs (all optional; sane defaults):
#   NOISE_ENABLED          master on/off switch (default false)
#   NOISE_PAD_TARGET       where to aim padding (default: the gateway Tailscale IP)
#   NOISE_PAD_PORT         UDP port on the target (default 9999; no listener needed)
#   NOISE_PAD_RATE_KBPS    target padding bitrate (default 64)
#   NOISE_PAD_PACKET_BYTES datagram payload size (default 1024)
#   NOISE_PAD_JITTER_MS    +/- timing jitter per packet (default 40)
#   NOISE_PAD_MODE         'constant' = always emit the target rate (default);
#                          'topup' = measure real egress each second and inject
#                          only the deficit, so the OBSERVED total stays flat
#                          whether idle or busy (stronger anti-correlation, no
#                          wasted bandwidth when you're already sending)
#   NOISE_PAD_PCT          set the target as a % of LINK CAPACITY (overrides
#                          NOISE_PAD_RATE_KBPS). % of capacity, NOT of current
#                          throughput  which would drop padding to zero when idle
#   NOISE_LINK_MBPS        link capacity for NOISE_PAD_PCT (default 100)
#   NOISE_PAD_IFACE        interface topup measures egress on (default en0)

set -uo pipefail

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
source "$SCRIPT_DIR/common.sh"
export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"

SELF="$SCRIPT_DIR/noise.sh"
NOISE_PID="/tmp/nullexit-noise.pid"
NOISE_LOG="$PROJECT_ROOT/output.log"

#  .env-driven config
NOISE_ENABLED=$(read_env_var "NOISE_ENABLED" | tr '[:upper:]' '[:lower:]')
[ -z "$NOISE_ENABLED" ] && NOISE_ENABLED="false"

PAD_TARGET=$(read_env_var "NOISE_PAD_TARGET")
[ -z "$PAD_TARGET" ] && PAD_TARGET="$(read_adguard_ip)"        # gateway Tailscale IP
PAD_PORT=$(read_env_var "NOISE_PAD_PORT");         [ -z "$PAD_PORT" ]  && PAD_PORT=9999
PAD_RATE=$(read_env_var "NOISE_PAD_RATE_KBPS");    [ -z "$PAD_RATE" ]  && PAD_RATE=64
PAD_BYTES=$(read_env_var "NOISE_PAD_PACKET_BYTES");[ -z "$PAD_BYTES" ] && PAD_BYTES=1024
PAD_JITTER=$(read_env_var "NOISE_PAD_JITTER_MS");  [ -z "$PAD_JITTER" ] && PAD_JITTER=40
# 'constant' = always emit the target rate. 'topup' = measure real egress each
# second and inject only the deficit, so the OBSERVED total rate stays flat
# whether you're idle or busy (the strong anti-correlation property).
PAD_MODE=$(read_env_var "NOISE_PAD_MODE" | tr '[:upper:]' '[:lower:]'); [ -z "$PAD_MODE" ] && PAD_MODE=constant
PAD_IFACE=$(read_env_var "NOISE_PAD_IFACE"); [ -z "$PAD_IFACE" ] && PAD_IFACE=en0
# Target rate: NOISE_PAD_PCT (a % of LINK CAPACITY  a stable reference, NOT of
# current throughput, which would collapse padding to zero when idle) overrides
# the absolute NOISE_PAD_RATE_KBPS when set.
PAD_PCT=$(read_env_var "NOISE_PAD_PCT")
PAD_LINK_MBPS=$(read_env_var "NOISE_LINK_MBPS"); [ -z "$PAD_LINK_MBPS" ] && PAD_LINK_MBPS=100
if [ -n "$PAD_PCT" ]; then
  PAD_RATE=$(awk "BEGIN{printf \"%.0f\", ($PAD_PCT/100.0)*$PAD_LINK_MBPS*1000}")
```

```
fi

log() { echo "[$(date '+%Y-%m-%d %H:%M:%S')] [noise] $*" >> "$NOISE_LOG" 2>/dev/null || true; }

require_enabled() {
  if [ "$NOISE_ENABLED" != "true" ]; then
    echo "noise: disabled (NOISE_ENABLED != true in .env)  refusing to emit traffic." >&2
    echo "       Set NOISE_ENABLED=true to arm cover-traffic padding." >&2
    exit 3
  fi
}

#  Randomness self-test (stdlib python3; no third-party deps)
# Exits 0 (PASS) / 1 (FAIL). Prints a human summary unless $2 == "quiet".
verify_random() {
  local nbytes="${1:-1048576}" mode="${2:-loud}"
  python3 - "$nbytes" "$mode" <<'PY'
import os, sys, math
from collections import Counter

n    = int(sys.argv[1])
mode = sys.argv[2] if len(sys.argv) > 2 else "loud"
buf  = os.urandom(n)
bits = n * 8

# 1) Shannon entropy over byte values (bits/byte, ideal = 8.0)
c = Counter(buf)
H = -sum((v / n) * math.log2(v / n) for v in c.values())

# 2) NIST monobit: 1-bit count should be ~half. z = (ones - bits/2)/sqrt(bits/4)
ones = sum(bin(b).count("1") for b in buf)
z    = (ones - bits / 2) / math.sqrt(bits / 4)

# 3) Chi-square over the 256 byte values (df = 255, mean 255, sd ~22.6)
exp  = n / 256.0
chi2 = sum((c.get(i, 0) - exp) ** 2 / exp for i in range(256))

# PASS windows chosen wide enough that true CSPRNG output passes ~always,
# yet a broken/patterned source (constant, low-entropy, biased) fails hard.
H_OK    = H >= 7.99 if n >= 1 << 16 else H >= 7.5
Z_OK    = abs(z) < 5.0
CHI_OK  = 185.0 < chi2 < 345.0
ok = H_OK and Z_OK and CHI_OK

if mode != "quiet":
    print(f"    sample        : {n} bytes ({bits} bits) from os.urandom (/dev/urandom CSPRNG)")
    print(f"    shannon H     : {H:.5f} bits/byte  (ideal 8.0)           [{'ok' if H_OK else 'FAIL'}]")
    print(f"    monobit z     : {z:+.3f}  (|z|<5, bias {ones/bits*100:.3f}% ones) [{'ok' if Z_OK else 'FAIL'}]"
     )
    print(f"    chi-square    : {chi2:.1f}  (df=255, expect ~255, window 185-345) [{'ok' if CHI_OK else 'FAIL
     '}]")
    print(f"    verdict       : {'PASS  statistically random' if ok else 'FAIL  not random enough'}")
sys.exit(0 if ok else 1)
PY
}

#  The cover-traffic loop (foreground; `start` backgrounds this)
pad_loop() {
  require_enabled
  # Own the PID file for BOTH run modes (manual `start` and launchd `__boot`), so
  # `status`/`stop` behave identically however padding was launched. Cleared on exit.
  echo $$ > "$NOISE_PID"
  trap 'rm -f "$NOISE_PID"' EXIT INT TERM
  # Gate: only ever emit noise we have just proven is statistically random.
  if ! verify_random 1048576 quiet; then
    echo "noise: randomness self-test FAILED  refusing to emit non-random padding." >&2
    log "pad aborted: randomness self-test failed"
    exit 1
  fi
  log "pad start  ${PAD_TARGET}:${PAD_PORT} mode=${PAD_MODE} rate=${PAD_RATE}kbps pkt=${PAD_BYTES}B jitter=${
     PAD_JITTER}ms iface=${PAD_IFACE}"
  python3 - "$PAD_TARGET" "$PAD_PORT" "$PAD_RATE" "$PAD_BYTES" "$PAD_JITTER" "$PAD_MODE" "$PAD_IFACE" <<'PY'
import os, sys, socket, time, random, subprocess

target     = sys.argv[1]
port       = int(sys.argv[2])
target_bps = float(sys.argv[3]) * 1000.0 / 8.0   # kbps  bytes/sec (target TOTAL on the wire)
pkt        = int(sys.argv[4])
jitter     = float(sys.argv[5]) / 1000.0         # ms  sec
mode       = sys.argv[6]                          # 'constant' | 'topup'
iface      = sys.argv[7]
```

```
s = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)


def obytes(ifc):
    """Cumulative output bytes for an interface  what an on-path observer sees
    leave the NIC. Returns None if the counter can't be read (fail toward MORE cover)."""
    try:
        out = subprocess.run(["netstat", "-ibn", "-I", ifc],
                             capture_output=True, text=True, timeout=2).stdout
        for line in out.splitlines():
            f = line.split()
            # the "<Link#N>" row carries cumulative counters; Obytes is column 9
            if len(f) >= 11 and f[0] == ifc and f[2].startswith("<Link"):
                return int(f[9])
    except Exception:
        pass
    return None


def emit(budget_bytes):
    """Send ~budget_bytes of fresh CSPRNG padding, spread across ~1s with jitter.
    Returns bytes actually sent (so topup can subtract its own contribution)."""
    n = max(0, int(budget_bytes // pkt))
    interval = 1.0 / n if n > 0 else 1.0
    sent = 0
    for _ in range(n):
        try:
            s.sendto(os.urandom(pkt), (target, port)); sent += pkt
        except OSError:
            pass  # dropped is fine  the bytes still left the NIC
        time.sleep(max(0.0, interval + random.uniform(-jitter, jitter)))
    return sent


try:
    prev = obytes(iface) if mode == "topup" else None
    pad_last = 0
    while True:
        t0 = time.time()
        if mode == "topup":
            now = obytes(iface)
            if prev is not None and now is not None:
                total = max(0, now - prev)        # all egress on iface this window
                real  = max(0, total - pad_last)  # minus our own last-window padding
                budget = max(0.0, target_bps - real)
            else:
                budget = target_bps               # counters unavailable  full target
            prev = now
        else:  # constant
            budget = target_bps
        pad_last = emit(budget)
        rem = 1.0 - (time.time() - t0)            # keep a ~1s accounting window
        if rem > 0:
            time.sleep(rem)
except KeyboardInterrupt:
    pass
PY
}

daemon_running() {
  [ -f "$NOISE_PID" ] || return 1
  local pid; pid="$(cat "$NOISE_PID" 2>/dev/null)"
  [ -n "$pid" ] && kill -0 "$pid" 2>/dev/null
}

cmd_start() {
  require_enabled
  if daemon_running; then
    echo "noise: padding daemon already running (pid $(cat "$NOISE_PID"))."
    exit 0
  fi
  echo "noise: verifying randomness before arming padding"
  verify_random 1048576 loud || { echo "noise: aborting  self-test failed." >&2; exit 1; }
  nohup "$SELF" pad >>"$NOISE_LOG" 2>&1 &
  local npid=$!
  sleep 0.3   # let pad_loop write its PID file so `status` is immediately accurate
  echo "noise: padding daemon armed (pid $npid)  ${PAD_TARGET}:${PAD_PORT}, ${PAD_RATE} kbps."
}

cmd_stop() {
```

```
  if daemon_running; then
    local pid; pid="$(cat "$NOISE_PID")"
    kill "$pid" 2>/dev/null && echo "noise: stopped padding daemon (pid $pid)."
    log "pad stop (pid $pid)"
  else
    echo "noise: no padding daemon running."
  fi
  rm -f "$NOISE_PID"
}

cmd_status() {
  echo "NOISE_ENABLED = $NOISE_ENABLED"
  if daemon_running; then
    echo "padding       : RUNNING (pid $(cat "$NOISE_PID"))"
    echo "target        : ${PAD_TARGET}:${PAD_PORT}"
    echo "mode          : ${PAD_MODE}$([ "$PAD_MODE" = topup ] && echo " (fills to target on ${PAD_IFACE};
     padding backs off as real traffic rises)")"
    echo "rate / packet : ${PAD_RATE} kbps${PAD_PCT:+ (=${PAD_PCT}% of ${PAD_LINK_MBPS} Mbps)} / ${PAD_BYTES} B
     , +/-${PAD_JITTER} ms jitter"
  else
    echo "padding       : stopped"
  fi
}

# launchd entry point (see launchd/com.nullexit.noise.plist). Governed by the
# SAME NOISE_ENABLED switch  NOT a new option. When disabled it exits 0 cleanly
# so KeepAlive never tight-loops; when enabled it runs the padding loop in the
# foreground so launchd supervises it directly. This is why persistence adds no
# extra knob: loading the plist once means "auto-start whenever NOISE_ENABLED=true".
cmd_boot() {
  if [ "$NOISE_ENABLED" != "true" ]; then
    log "boot: NOISE_ENABLED != true  idle (padding not started)"
    exit 0
  fi
  log "boot: NOISE_ENABLED=true  starting padding loop under launchd"
  pad_loop
}

case "${1:-}" in
  verify)      verify_random "${2:-1048576}" loud ;;
  start)       cmd_start ;;
  pad)         pad_loop ;;
  __boot)      cmd_boot ;;
  stop)        cmd_stop ;;
  status)      cmd_status ;;
  --help|-h|"") sed -n '2,59p' "$0" ;;
  *) echo "Unknown argument: $1 (try --help)" >&2; exit 2 ;;
esac
```

45 scripts/pf.conf

```
# scripts/pf.conf - Packet Filter configuration for nullexit killswitch
# Dynamically parameterized by pfctl macro override

# ext_if must be passed via command line, e.g. -D ext_if=en0
# If not passed, default to en0
ext_if = "en0"
ts_if = "utun+"

# Tables populated dynamically via pfctl
table <vpn_endpoints> persist
table <derp_relays> persist
# FaceTime split-route (opt-in): Apple media/relay /16s that egress DIRECT via the
# physical gateway. EMPTY unless FACETIME_SPLIT_ROUTE=true, so the kill-switch
# posture is unchanged by default. Populated by add_apple_split_routes().
table <apple_direct> persist

# Core Rules
set skip on lo0                # Ensure local loopback is never blocked
# TCP MSS Clamping: prevents fragmentation on host TCP flows (like Tailscale control plane)
# by keeping the host in sync with the container's double-tunnel (Tailscale+WARP) overhead.
# NOTE: this 1120 integer is rewritten at kill-switch enable time by scripts/common.sh's
# enable_killswitch() from GATEWAY_MSS in .env  change GATEWAY_MSS in .env, not this line.
scrub out all max-mss 1120
pass quick on $ts_if all       # Allow all Tailscale traffic (essential to prevent SSH lockout)

```

```
# Default deny outbound WAN traffic
block out on $ext_if all

# Block inbound Screen Sharing (VNC) and SSH from local Wi-Fi, forcing them to use Tailscale
block in on $ext_if proto tcp from any to any port { 22, 5900 }

# Exceptions for tunnel handshakes and coordination
pass out quick on $ext_if proto udp to <vpn_endpoints> port 2408
pass out quick on $ext_if proto udp to <derp_relays>
pass out quick on $ext_if proto tcp to 192.200.0.0/24 port 443

# FaceTime split-route (opt-in, table empty unless FACETIME_SPLIT_ROUTE=true):
# let real-time media/relay traffic to Apple's /16s leave DIRECT (real IP) so
# FaceTime can hole-punch instead of dying behind WARP's symmetric NAT.
pass out quick on $ext_if proto { tcp udp } to <apple_direct>

# Allow local LAN traffic (AirDrop, local Tailscale P2P, etc)
pass out quick on $ext_if proto tcp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
pass out quick on $ext_if proto udp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
pass out quick on $ext_if proto icmp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }

# Allow DHCP outbound (client port 68 to server port 67) so we can obtain/renew IP address
pass out quick on $ext_if proto udp from any port 68 to any port 67
```

46 scripts/pihole-mode.sh

```
#!/bin/bash
# scripts/pihole-mode.sh  LAN Pi-hole mode (opt-in)
#
# nullexit is already a Pi-hole for MESH devices: AdGuard Home sinkholes DNS for
# everything that routes through the gateway over Tailscale. This mode extends
# that to NON-Tailscale devices on the same LAN  a smart TV, a guest laptop, an
# IoT gadget  the classic "point your DNS at this box" Pi-hole. It binds a DNS
# forwarder on the LAN interface that relays to AdGuard, so any device that sets
# its DNS server to this Mac's LAN IP gets the same ad/tracker/threat filtering.
#
# Mechanism: reuse the tested dns-proxy.py, but LISTEN_ADDR = the LAN IP instead
# of loopback, forwarding to AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (the Docker
# port mapping). No kill-switch changes are needed  pf.conf has no inbound-53
# block, and DNS responses to RFC1918 LAN clients are already permitted.
#
#  NETWORK REALITY: on an AP-isolated network (WPA2-Enterprise / hotel), other
# devices CANNOT reach this Mac at all  the same isolation that forces phones
# onto DERP. This mode only does anything useful on a TRUSTED, non-isolated
# network (home / your own AP). It is report-clear here but won't serve remote
# clients until you're on such a network. It also filters DNS only  it does NOT
# route those devices' traffic through WARP.
#
# Usage:
#   bash scripts/pihole-mode.sh enable    # start the LAN DNS forwarder (needs PIHOLE_LAN_MODE=true)
#   bash scripts/pihole-mode.sh disable   # stop it
#   bash scripts/pihole-mode.sh status    # running? listen addr, AdGuard reachability
#   bash scripts/pihole-mode.sh test      # prove it resolves AND sinkholes, locally
#
# .env:
#   PIHOLE_LAN_MODE     master on/off (default false)
#   PIHOLE_LAN_LISTEN   LAN address to bind (blank = this Mac's en0 IP)
#   DNS_PROXY_PORT      AdGuard's host-published DNS port (default 5354)

set -uo pipefail

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
source "$SCRIPT_DIR/common.sh"
export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"

DNS_PROXY="$SCRIPT_DIR/dns-proxy.py"
PID_FILE="/tmp/nullexit-pihole.pid"
LOG_FILE="$PROJECT_ROOT/output.log"

PIHOLE_ON=$(read_env_var "PIHOLE_LAN_MODE" | tr '[:upper:]' '[:lower:]'); [ -z "$PIHOLE_ON" ] && PIHOLE_ON="
    false"

# AdGuard's DNS host port is PROCEDURAL (randomized per boot  see Procedural
# Ports), so discover it; never assume 5354. Docker is the source of truth, then
```

```
# the ports file toggle.sh writes, then the compose default.
resolve_adguard_port() {
  local p
  p=$(docker compose port warp 5335 2>/dev/null | grep -oE '[0-9]+$' | head -1)
  [ -n "$p" ] && { echo "$p"; return; }
  p=$(grep -oE 'DNS_PROXY_PORT=[0-9]+' /tmp/nullexit-ports.env 2>/dev/null | grep -oE '[0-9]+$' | head -1)
  [ -n "$p" ] && { echo "$p"; return; }
  echo 5354
}
ADGUARD_PORT=$(resolve_adguard_port)

lan_ip() {
  local a; a=$(read_env_var "PIHOLE_LAN_LISTEN")
  [ -n "$a" ] && { echo "$a"; return; }
  ipconfig getifaddr en0 2>/dev/null || echo ""
}

log() { echo "[$(date '+%Y-%m-%d %H:%M:%S')] [pihole-lan] $*" >> "$LOG_FILE" 2>/dev/null || true; }

running() {
  [ -f "$PID_FILE" ] || return 1
  local p; p=$(cat "$PID_FILE" 2>/dev/null)
  [ -n "$p" ] && kill -0 "$p" 2>/dev/null
}

adguard_reachable() {
  # Probe over TCP: Colima user-mode net forwards UDP-to-container unreliably (the
  # very reason dns-proxy.py relays over TCP), so a UDP probe false-negatives even
  # when AdGuard is healthy. TCP is exactly the path the forwarder uses.
  dig +tcp +short +tries=1 +time=2 -p "$ADGUARD_PORT" @127.0.0.1 example.com >/dev/null 2>&1
}

cmd_enable() {
  if [ "$PIHOLE_ON" != "true" ]; then
    echo "pihole-lan: disabled (PIHOLE_LAN_MODE != true in .env)  refusing to serve LAN DNS." >&2
    exit 3
  fi
  if running; then echo "pihole-lan: already running (pid $(cat "$PID_FILE"))."; exit 0; fi
  local addr; addr=$(lan_ip)
  if [ -z "$addr" ]; then echo "pihole-lan: could not determine a LAN IP (en0 down?)." >&2; exit 1; fi
  if ! adguard_reachable; then
    echo "pihole-lan: AdGuard not reachable at 127.0.0.1:$ADGUARD_PORT  is the gateway up?" >&2
    exit 1
  fi
  echo "pihole-lan: starting LAN DNS forwarder on ${addr}:53  AdGuard 127.0.0.1:${ADGUARD_PORT}"
  # :53 needs root. The shipped sudoers NOPASSWD rule only covers the literal
  # `python3 dns-proxy.py` command  no bash -c wrapper, no env prefix  so
  # config is dropped in a file dns-proxy.py reads itself, not via env vars
  # that `sudo -n` would strip. See dns-proxy.py's own comment for why.
  {
    echo "LISTEN_ADDR=$addr"
    echo "LISTEN_PORT=53"
    echo "TARGET_HOST=127.0.0.1"
    echo "TARGET_PORT=$ADGUARD_PORT"
  } > /tmp/nullexit-dns-proxy.conf
  chmod 600 /tmp/nullexit-dns-proxy.conf 2>/dev/null || true
  nohup sudo -n "$(command -v python3)" "$DNS_PROXY" >>"$LOG_FILE" 2>&1 &
  disown "$!" 2>/dev/null || true
  echo "$!" > "$PID_FILE"
  sleep 0.4
  if running; then
    log "enabled on ${addr}:53  127.0.0.1:${ADGUARD_PORT}"
    echo "pihole-lan: running (pid $(cat "$PID_FILE")). Point a LAN device's DNS at ${addr} to filter it."
  else
    echo "pihole-lan: listener did not stay up  check $LOG_FILE (port 53 already bound?)." >&2; exit 1
  fi
}

cmd_disable() {
  if running; then
    local p; p=$(cat "$PID_FILE")
    sudo -n kill "$p" 2>/dev/null || kill "$p" 2>/dev/null
    log "disabled (pid $p)"; echo "pihole-lan: stopped (pid $p)."
  else
    echo "pihole-lan: not running."
  fi
  sudo -n rm -f "$PID_FILE" 2>/dev/null || rm -f "$PID_FILE" 2>/dev/null || true
  rm -f /tmp/nullexit-dns-proxy.conf 2>/dev/null || true
}

cmd_status() {
```

```
  echo "PIHOLE_LAN_MODE = $PIHOLE_ON"
  echo "AdGuard (127.0.0.1:$ADGUARD_PORT): $(adguard_reachable && echo reachable || echo UNREACHABLE)"
  if running; then
    echo "forwarder      : RUNNING (pid $(cat "$PID_FILE")) on $(lan_ip):53"
  else
    echo "forwarder      : stopped"
  fi
}

cmd_test() {
  local addr; addr=$(lan_ip)
  if ! running; then echo "pihole-lan: not running  'enable' first." >&2; exit 1; fi
  echo "Testing LAN Pi-hole at ${addr}:53 (querying locally)"
  local good bad
  good=$(dig +short +tries=1 +time=3 @"$addr" example.com 2>/dev/null | head -1)
  bad=$(dig +short +tries=1 +time=3 @"$addr" doubleclick.net 2>/dev/null | head -1)
  echo "  resolve example.com      ${good:-<none>}"
  echo "  filter  doubleclick.net  ${bad:-<blocked/empty>}"
  if [ -n "$good" ] && { [ -z "$bad" ] || [ "$bad" = "0.0.0.0" ]; }; then
    echo "  RESULT: PASS  resolves clean domains and sinkholes blocked ones."
  else
    echo "  RESULT: check  clean resolve=${good:+yes}, sinkhole=${bad:-empty}. (If both empty, AdGuard/gateway
     may be down.)"
  fi
}

case "${1:-}" in
  enable)  cmd_enable ;;
  disable) cmd_disable ;;
  status)  cmd_status ;;
  test)    cmd_test ;;
  --help|-h|"") sed -n '2,40p' "$0" ;;
  *) echo "Unknown argument: $1 (try --help)" >&2; exit 2 ;;
esac
```

47 scripts/profiles.sh

```
#!/bin/bash
# scripts/profiles.sh  named config profiles that tame the .env flag sprawl.
#
# CONFIG-ONLY BY DESIGN: this tool ONLY reads/writes .env. It NEVER touches the
# live gateway  no routing, no PF, no containers, no restart  so it cannot break
# your running network. Changes take effect only on your NEXT ./toggle.sh --restart.
#
# A profile is a coherent, known-good bundle of the intent-level switches, so you
# pick ONE intent instead of reasoning about ~7 interacting flags. KILL_SWITCH is
# forced true in every profile  the fail-closed invariant is never negotiable.
#
# Usage:
#   bash scripts/profiles.sh show            # current config, grouped + explained (read-only)
#   bash scripts/profiles.sh list            # available profiles and what each sets (read-only)
#   bash scripts/profiles.sh preview <name>  # exactly what 'apply' would change (read-only)
#   bash scripts/profiles.sh apply <name>    # write the profile to .env (backs up; never restarts)
#
# Profiles: campus (AP-isolated enterprise, conservative)  home (trusted, fast)
#           hardened (max privacy: cover traffic + Tor bridges, no real-IP exposure)

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
cd "$PROJECT_ROOT" || exit 1
case "${1:-}" in --help|-h) sed -n '2,25p' "$0"; exit 0 ;; esac

python3 - "$@" <<'PY'
import sys, os, shutil, time

ENV = os.environ.get("PROFILES_ENV_FILE", ".env")

# Intent-level switches each profile manages. KILL_SWITCH is a forced invariant.
PROFILES = {
  "campus":   {"KILL_SWITCH":"true","TAILSCALE_ALLOW_LAN_P2P":"false","FACETIME_SPLIT_ROUTE":"false",
               "PIHOLE_LAN_MODE":"false","NOISE_ENABLED":"false","TOR_USE_BRIDGES":"false","WARP_FAIL_THRESHOLD
    ":"6"},
  "home":     {"KILL_SWITCH":"true","TAILSCALE_ALLOW_LAN_P2P":"true","FACETIME_SPLIT_ROUTE":"true",
               "PIHOLE_LAN_MODE":"true","NOISE_ENABLED":"false","TOR_USE_BRIDGES":"false","WARP_FAIL_THRESHOLD"
    :"6"},
  "hardened": {"KILL_SWITCH":"true","TAILSCALE_ALLOW_LAN_P2P":"false","FACETIME_SPLIT_ROUTE":"false",
```

```
                "PIHOLE_LAN_MODE":"false","NOISE_ENABLED":"true","NOISE_PAD_MODE":"topup",
                "TOR_USE_BRIDGES":"true","WARP_FAIL_THRESHOLD":"3"},
}
DESC = {
  "campus":"AP-isolated enterprise/campus: DERP (no SNAT poisoning), everything conservative, no real-IP
     exposure.",
  "home":"Trusted network: direct LAN P2P (fast), FaceTime split-route + LAN Pi-hole on, no cover-traffic
     battery cost.",
  "hardened":"Max privacy on a hostile net: cover traffic (topup) + Tor bridges, fail-fast WARP, never expose
     the real IP.",
}
META = {
  "KILL_SWITCH":"PF fail-closed kill-switch (forced true  never negotiable)",
  "TAILSCALE_ALLOW_LAN_P2P":"direct LAN P2P vs forced DERP relay",
  "FACETIME_SPLIT_ROUTE":"route Apple /16s direct (fixes FaceTime; exposes real IP for them)",
  "PIHOLE_LAN_MODE":"serve AdGuard DNS to non-mesh LAN devices",
  "NOISE_ENABLED":"cover-traffic padding",
  "NOISE_PAD_MODE":"padding rate mode (constant / topup)",
  "TOR_USE_BRIDGES":"Tor via obfs4 bridges only",
  "WARP_FAIL_THRESHOLD":"consecutive warp=off polls before auto-shutdown",
}

def read_env():
    if not os.path.exists(ENV):
        print(f"profiles: {ENV} not found", file=sys.stderr); sys.exit(1)
    return open(ENV, encoding="utf-8").read().splitlines()

def get_val(lines, key):
    for l in lines:
        s = l.strip()
        if s.startswith(key + "="):
            return s.split("=", 1)[1]
    return None

def which_profile(lines):
    """Name the profile that matches the current .env, or None."""
    for name, flags in PROFILES.items():
        if all(get_val(lines, k) == v for k, v in flags.items()):
            return name
    return None

def cmd_show():
    lines = read_env()
    active = which_profile(lines)
    print(f"\n=== nullexit config (current){'   profile: ' + active if active else '   custom (no exact profile
      match)'} ===\n")
    for k, d in META.items():
        v = get_val(lines, k)
        print(f"  {k:26} = {str(v):9}  {d}")
    print("\n  Read-only. 'list' to see profiles, 'apply <name>' for a known-good bundle.")

def cmd_list():
    lines = read_env()
    print()
    for name, flags in PROFILES.items():
        print(f"   {name}  {DESC[name]}")
        for k, v in flags.items():
            cur = get_val(lines, k)
            print(f"      {k:26}  {v}{'' if cur == v else '   (now: ' + str(cur) + ')'}")
        print()

def diff_for(name):
    lines = read_env()
    return [(k, get_val(lines, k), v) for k, v in PROFILES[name].items() if get_val(lines, k) != v]

def cmd_preview(name):
    if name not in PROFILES:
        print(f"unknown profile '{name}' (try: {', '.join(PROFILES)})", file=sys.stderr); sys.exit(2)
    ch = diff_for(name)
    print(f"\nprofile '{name}'  {DESC[name]}\n")
    if not ch:
        print("  already applied  no changes."); return
    print("  would change:")
    for k, cur, v in ch:
        print(f"    {k:26} {cur}  {v}")
    print("\n  Preview only  nothing written, network untouched.")

def set_key(lines, key, val):
    for i, l in enumerate(lines):
        if l.strip().startswith(key + "="):
            lead = l[:len(l) - len(l.lstrip())]
```

```
            lines[i] = f"{lead}{key}={val}"
            return
    lines.append(f"{key}={val}")

def cmd_apply(name):
    if name not in PROFILES:
        print(f"unknown profile '{name}' (try: {', '.join(PROFILES)})", file=sys.stderr); sys.exit(2)
    ch = diff_for(name)
    if not ch:
        print(f"profile '{name}' already applied  no changes."); return
    lines = read_env()
    bak = f"{ENV}.bak.{time.strftime('%Y%m%d%H%M%S')}"
    shutil.copy2(ENV, bak)
    for k, v in PROFILES[name].items():
        set_key(lines, k, v)
    with open(ENV, "w", encoding="utf-8") as f:
        f.write("\n".join(lines) + "\n")
    print(f"\napplied '{name}' to {ENV}  (backup: {bak})\n")
    for k, cur, v in ch:
        print(f"    {k:26} {cur}  {v}")
    print("\n    .env only  your LIVE gateway is UNCHANGED. Apply with:  ./toggle.sh --restart")

cmd = sys.argv[1] if len(sys.argv) > 1 else "show"
arg = sys.argv[2] if len(sys.argv) > 2 else ""
{"show": cmd_show, "list": cmd_list,
 "preview": (lambda: cmd_preview(arg)), "apply": (lambda: cmd_apply(arg))}.get(cmd, cmd_show)()
PY
```

48 scripts/reinstall-host-tailscale.sh

```
#!/usr/bin/env bash
#
# scripts/reinstall-host-tailscale.sh
#
# Complete uninstall + reinstall of the host-side Homebrew Tailscale.
#
# Use this when the brew LaunchAgent is wedged, Login Items entries are
# duplicated, "Tailscale is stopped" persists despite
# `brew services restart tailscale`, or you simply want a clean baseline
# before re-engaging the gateway as exit node.
#
# Idempotent. Safe to re-run. Use --dry to preview every step with zero
# changes. Use --yes to skip the per-step confirmation prompts.
#
# This script ONLY touches the HOST-side Tailscale (the one installed by
# Homebrew `brew install tailscale` and managed by launchd as
# homebrew.mxcl.tailscale). The nullexit gateway container's tailscaled
# (which lives inside the Colima VM at 100.100.21.8 and offers the exit
# node to this host) is unaffected.

set -u

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"

DRY=0
YES=0

for arg in "$@"; do
  case "$arg" in
    --dry)  DRY=1 ;;
    --yes)  YES=1 ;;
    --help|-h)      # NOTE: heredoc is UNQUOTED so $SCRIPT_DIR expands in the body.  If you
      # add new $-style placeholders here, they'll expand the same way.
      cat <<USAGE
Usage: bash "$SCRIPT_DIR/reinstall-host-tailscale.sh" [--dry] [--yes]

  --dry  Print every step, make zero changes. Use to audit before
         committing.
  --yes  Auto-confirm the prompts at every destructive step. Useful in
         CI / fully-scripted recovery; do not use until you've read
         the script.

After this script completes cleanly, you still must (MANUALLY):

  1. Open System Settings  Privacy & Security  Local Network and
     enable Tailscale. macOS will prompt automatically the first time
```

```
     the fresh install of `tailscale` is invoked.

  2. Re-engage the tailnet + exit node:
       sudo tailscale up --ssh=true --accept-dns=false \
         --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
         --exit-node-allow-lan-access=true

  3. Re-run ./toggle.sh from the project root.

  4. Final verification:
       curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
    expected: warp=on, ip=<Cloudflare>, colo=YYZ/YUL
USAGE
      exit 0
      ;;
    *)
      echo "Unknown flag: $arg" >&2
      exit 2
      ;;
  esac
done

# ---------- helpers ----------

confirm() {
  if [ "$YES" = 1 ] || [ "$DRY" = 1 ]; then
    return 0
  fi
  printf "      %s  [y/N] " "$1"
  read -r ans
  case "$ans" in
    y|Y|yes|YES) return 0 ;;
    *)
      echo "     aborted by user"
      exit 1
      ;;
  esac
}

header() {
  printf '\n %s\n' "$1"
}

# with_timeout <seconds> <cmd...>
# Run a command in the background and kill it if it exceeds <seconds>.
# Returns the command's exit code if it finished in time, 124 (matching
# GNU `timeout` convention) if it timed out. macOS doesn't ship GNU
# coreutils `timeout`, so this is the bash-3.2-friendly replacement.
# stdout + stderr are discarded (we only care about exit code); redirect
# or pipe externally if you need the output.
with_timeout() {
  local timeout_s="${1:-30}"
  shift
  "$@" >/dev/null 2>&1 &
  local cmdpid=$!
  local elapsed=0
  while kill -0 "$cmdpid" 2>/dev/null; do
    if [ "$elapsed" -ge "$timeout_s" ]; then
      kill -9 "$cmdpid" 2>/dev/null
      wait "$cmdpid" 2>/dev/null
      return 124
    fi
    sleep 1
    elapsed=$((elapsed+1))
  done
  wait "$cmdpid" 2>/dev/null
  return $?
}

# ---------- 0. Pre-flight ----------
header "Pre-flight  current host tailscale state"
echo "    brew:                  $(command -v brew >/dev/null 2>&1 && brew --version | head -1 || echo 'not on
      PATH')"
echo "    tailscaled running:    $(pgrep -lf tailscaled 2>/dev/null | wc -l | tr -d ' ') process(es)"
echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{print $2, $3}' | tr
     '\n' ' ' | sed 's/ $//')"
echo "    LaunchAgent plists (under ~/Library/LaunchAgents/):"
ls -1 ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/^/      /' || echo "      (none)"

# Detect Tailscale's macOS Network Extension (residual from prior App Store
# Tailscale.app installs). Lives under /Library/Apple/SystemExtensions/ and
```

```
# is managed by `systemextensionsctl`  NOT launchd. If loaded, brew
# tailscaled will NOT engage the Network Extension framework slot because
# the SE still has it claimed. Symptom: 'sudo tailscale status' returns
# "Tailscale is stopped" even when the daemon is alive and listening.
se_found=0
if command -v systemextensionsctl >/dev/null 2>&1; then
  # Grab any line matching tailscale and network-extension
  se_line=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale\..*network-extension' | head -n 1)
  if [ -n "$se_line" ]; then
    se_found=1
    # Extract the teamID (10-char uppercase/numeric word)
    se_teamID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^[A-Z0-9]{10}$' | head -n 1)
    # Extract the bundleID (starts with io.tailscale or com.tailscale)
    se_bundleID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^(io|com)\.tailscale\.' | head -n 1 | sed 's/(.*//')

    # If teamID is not found, fallback to '-'
    if [ -z "$se_teamID" ]; then
      se_teamID="-"
    fi
  fi
fi

if [ "$se_found" = 1 ] && [ -n "$se_bundleID" ]; then
  echo "     Tailscale Network Extension System Extension is loaded (residual from App Store Tailscale.app install)"
  echo "      Team ID: $se_teamID, Bundle ID: $se_bundleID"
  echo "      brew tailscaled cannot engage the NE slot while this SE has it claimed."
  if [ "$DRY" = 1 ]; then
    echo "    [dry] (would prompt) sudo systemextensionsctl uninstall $se_teamID $se_bundleID"
  else
    confirm "Run 'sudo systemextensionsctl uninstall $se_teamID $se_bundleID'? (frees the NE slot for brew tailscale; non-fatal if it errors)"
    sudo systemextensionsctl uninstall "$se_teamID" "$se_bundleID" 2>&1 \
      || echo "    ! uninstall returned non-zero (may require reboot, non-fatal)"
  fi
else
  echo "     no Tailscale System Extension loaded"
fi

# Pattern used by the Step-1 bootout loops, in both dry + real branches.
# Widened from a bare /tailscale/ so the same loops also pick up any
# nullexit-labeled LaunchAgents (e.g., com.nullexit.wake-recovery, the
# caffeinate + post-wake recovery hook). The wake-recovery LaunchAgent
# is re-installed by ./toggle.sh / setup.sh on re-engage, so wiping it
# here is non-fatal and is part of restoring a clean baseline.
TAILSCALE_LABEL_RE='tailscale'
NULLEXIT_LABEL_RE='nullexit'

# ---------- 1. brew services stop ----------
header "Step 1  Stop the brew-managed service + unload LaunchAgents"
# Detect whether the brew-managed service is registered as $USER (gui/$UID
# domain) or as root (system domain). NOPASSWD sudoers can leave a stale
# root-owned registration after a prior 'sudo brew services ...' round;
# 'brew services stop' without sudo refuses to touch a root-owned plist
# (Error: Service 'tailscale' is started as `root`).
brew_status_line="$(brew services list 2>/dev/null | awk '$1=="tailscale"')"
brew_status_user="$(echo "$brew_status_line" | awk '{print $3}')"
if [ -n "$brew_status_line" ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew services stop tailscale (user=$brew_status_user)"
  else
    if [ "$brew_status_user" = "root" ]; then
      confirm "Run 'sudo brew services stop tailscale' (service is registered as root)?"
      sudo brew services stop tailscale 2>&1 || echo "    (sudo brew services stop returned non-zero  non-fatal)"
    else
      echo "     brew services stop tailscale (user=$brew_status_user, no sudo needed)"
      brew services stop tailscale 2>&1 || echo "    (brew reported it wasn't actually running)"
    fi
  fi
  echo "     unload any tailscale-related LaunchAgents (prevents launchd from respawning after kill)"
  if [ "$DRY" = 1 ]; then
    echo "      [dry] for each 'tailscale|nullexit'-matched label in gui/$UID domain:"
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "            [dry] launchctl bootout gui/$UID/$label"
    done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
    echo "      [dry] for each 'tailscale|nullexit'-matched label in system domain (root-owned brew plist):"
```

```
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "             [dry] (would prompt) sudo launchctl bootout system/$label"
    done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
     NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
    if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
      echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macos"
    fi
    if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
      echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macsys"
    fi
  else
    # Bootout every gui-domain LaunchAgent whose label matches 'tailscale'
    # OR 'nullexit'. The nullexit arm catches com.nullexit.wake-recovery,
    # which holds caffeinate/watcher.sh alive across a daemon reset  not
    # desired while the script is establishing a fresh baseline. Awake-
    # recovery is re-installed by ./toggle.sh / setup.sh on re-engage, so
    # wiping here is non-fatal. The pattern stays narrower than a bare
    # /tail/ or /null/ so unrelated Apple services are not unloaded.
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "       launchctl bootout gui/$UID/$label"
      launchctl bootout "gui/$UID/$label" 2>/dev/null || true
    done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 &&
      ($3 ~ ts || $3 ~ nx) {print $3}')
    # Bootout every system-domain LaunchDaemon/Agent whose label matches
    # 'tailscale' OR 'nullexit'. This catches the root-owned brew plist left
    # behind by prior 'sudo brew services ...' invocations  that was the
    # case that left tailscaled PID 47447  57734 respawning between sudo
    # pkill runs. Defensive nullexit catch in case the user previously
    # `sudo ln -s`'d the wake-recovery plist into /Library/LaunchDaemons.
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      confirm "Run 'sudo launchctl bootout system/$label'?"
      sudo launchctl bootout "system/$label" 2>/dev/null || true
    done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
     NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
    # Legacy Tailscale.app system-domain entries. Each one is checked and
    # confirmed separately so we don't sudo-promp or sudo-bootout for nothing.
    if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
      confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macos'?"
      sudo launchctl bootout system/com.tailscale.ipn.macos 2>/dev/null || true
    fi
    if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
      confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macsys'?"
      sudo launchctl bootout system/com.tailscale.ipn.macsys 2>/dev/null || true
    fi
  fi
  echo "     waiting up to 10s for tailscaled to exit"
  for i in 1 2 3 4 5 6 7 8 9 10; do
    if [ "$DRY" = 1 ]; then
      echo "    [dry] sleep 1; loop check would terminate here"
      break
    fi
    sleep 1
    if ! pgrep -lf tailscaled >/dev/null 2>&1; then
      echo "     tailscaled exited in ${i}s"
      break
    fi
    [ "$i" = 10 ] && echo "     tailscaled still alive after 10s  Step 2 will deal with it"
  done
else
  echo "     brew reports no tailscale service loaded"
fi

# ---------- 2. Kill any orphan tailscaled ----------
header "Step 2  Kill any orphan tailscaled process(es) (may escalate to sudo)"
if ! pgrep -lf tailscaled >/dev/null 2>&1; then
  echo "     no orphans"
else
  echo "    Stray process(es):"
  pgrep -lf tailscaled | sed 's/^/      /'
  echo "     no-sudo: pkill tailscaled"
  if [ "$DRY" = 1 ]; then
    echo "    [dry] pkill tailscaled  on failure: sudo pkill -9 tailscaled"
  else
    if pkill tailscaled 2>/dev/null; then
      echo "     no-sudo pkill succeeded"
    else
      echo "    ! no-sudo pkill insufficient; cannot escalate yet"
    fi
```

```
    sleep 1
  fi

  # Recheck; if still alive, offer sudo escalation
  if { [ "$DRY" = 0 ] && pgrep -lf tailscaled >/dev/null 2>&1; } || [ "$DRY" = 1 ]; then
    if [ "$DRY" = 1 ]; then
      echo "    [dry] sudo pkill -9 tailscaled  (escalation would fire here)"
    else
      confirm "Run 'sudo pkill -9 tailscaled' to escalate?"
      sudo pkill -9 tailscaled 2>&1
      sleep 1
    fi
  fi

  if [ "$DRY" = 0 ]; then
    if pgrep -lf tailscaled >/dev/null 2>&1; then
      echo "     Failed to kill stragglers  bailing out"
      pgrep -lf tailscaled | sed 's/^/      /'
      exit 3
    fi
    echo "     all tailscaled processes gone"
  fi
fi

# ---------- 3. brew uninstall ----------
header "Step 3  brew uninstall tailscale"
if ! brew list tailscale >/dev/null 2>&1; then
  echo "     not currently installed via brew  skip"
else
  # Detect whether the keg directory is owner=$USER or owner=root. If the
  # keg is root-owned (typically a residual of some prior 'sudo brew ...'
  # cycle that did perms fixups), 'brew uninstall' without sudo refuses
  # to remove the files and prints: 'Error: Could not remove tailscale
  # keg! Do so manually: sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>'.
  # Detect that and escalate to a sudoed uninstall, with a manual rm -rf
  # fallback if `sudo brew uninstall` itself errors out (e.g., partial
  # state where the keg dir is gone but brew still thinks it's installed).
  # Detect keg state. Treat `(dir absent)` AND `(dir present + empty)` as the
  # same `<unknown>` bucket so both escalate to sudo. The empty-parent case
  # is what you get after a manual `sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>`
  # from inside: brew's parent dir is left present-but-empty and brew still
  # thinks the package is installed.
  keg_dir="/opt/homebrew/Cellar/tailscale"
  if [ -d "$keg_dir" ] && [ -n "$(ls -A "$keg_dir" 2>/dev/null)" ]; then
    keg_owner="$(stat -f '%Su' "$keg_dir" 2>/dev/null || echo '<unknown>')"
  else
    keg_owner="<unknown>"
  fi
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew uninstall tailscale (keg_owner=$keg_owner)"
  else
    # 'root' AND '<unknown>' both require sudo escalation. The '<unknown>'
    # case is the partial-state path: keg dir was already removed (manually
    # or by a prior failed uninstall) but brew's internal state still says
    # installed. Without this branch the script would fall through to a
    # non-sudo brew uninstall that fails again with the same 'Could not
    # remove tailscale keg!' error.
    if [ "$keg_owner" = "root" ] || [ "$keg_owner" = "<unknown>" ]; then
      confirm "Run 'sudo brew uninstall tailscale' (keg=$keg_owner)?"
      sudo brew uninstall tailscale 2>&1 || {
        echo "    ! sudo brew uninstall failed; offering manual rm -rf fallback"
        # Guard the glob so an empty/absent dir doesn't literal-expand the
        # '*' and confuse the user with a sudo error about a non-existent path.
        if [ -d /opt/homebrew/Cellar/tailscale ] && [ -n "$(ls -A /opt/homebrew/Cellar/tailscale 2>/dev/null)" ]; then
          confirm "Run 'sudo rm -rf /opt/homebrew/Cellar/tailscale/*'?"
          sudo rm -rf /opt/homebrew/Cellar/tailscale/* 2>&1
        else
          echo "     keg dir already empty or absent  nothing more to rm"
        fi
      }
    else
      confirm "Run 'brew uninstall tailscale' (removes LaunchAgent plist + linked symlinks)?"
      brew uninstall tailscale 2>&1
    fi
  fi
fi

# ---------- 4. Wipe state directories ----------
header "Step 4  Wipe stale state"
# User-owned
```

```
declare -a USER_PATHS
USER_PATHS=(
  "$HOME/.local/share/tailscale"
  "$HOME/.local/run/tailscaled.socket"
  "$HOME/.local/run/tailscaled.state"
  "$HOME/.local/state/tailscale"
  "$HOME/Library/Application Support/Tailscale"
  "$HOME/Library/Caches/com.tailscale.ipn.macos"
  "$HOME/Library/Caches/Tailscale"
  "$HOME/Library/Logs/Tailscale"
)
echo "    User-owned paths (no sudo):"
for p in "${USER_PATHS[@]}"; do
  if [ -e "$p" ]; then
    if [ "$DRY" = 1 ]; then
      echo "      [dry] rm -rf '$p'"
    else
      rm -rf "$p"
      echo "       rm -rf '$p'"
    fi
  fi
done

# System
declare -a SUDO_PATHS
SUDO_PATHS=(
  "/var/lib/tailscale"
  "/var/cache/tailscale"
)
echo "    System paths under /var (will prompt for sudo if any are present):"
eligible_sudo=0
for p in "${SUDO_PATHS[@]}"; do
  if [ -e "$p" ]; then
    eligible_sudo=1
    if [ "$DRY" = 1 ]; then
      echo "      [dry] sudo rm -rf '$p'"
    else
      if confirm "Run 'sudo rm -rf $p'?"; then
        sudo rm -rf "$p" && echo "       rm -rf '$p'"
      fi
    fi
  fi
done
[ "$eligible_sudo" = 0 ] && echo "      (none of /var/lib/tailscale, /var/cache/tailscale are present)"

# ---------- 5. Wipe stale LaunchAgent plists ----------
header "Step 5  Wipe stale LaunchAgent plist files"
declare -a PLIST_FILES
PLIST_FILES=(
  "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
  "$HOME/Library/LaunchAgents/com.tailscale.ipn.macos.plist"
  "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist.bak"
  "$HOME/Library/LaunchAgents/com.nullexit.wake-recovery.plist"
)
for f in "${PLIST_FILES[@]}"; do
  if [ -e "$f" ]; then
    if [ "$DRY" = 1 ]; then
      echo "    [dry] rm -f '$f'"
    else
      rm -f "$f"
      echo "     rm -f '$f'"
    fi
  fi
done
echo "    Remaining tailscale|nullexit plists (should be empty):"
ls -1 ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/^/      /' || echo "      (none)"

# ---------- 5.5. Drain Background-Items (Login Items) ----------
header "Step 5.5  Drain Background-Items (Login Items)"
# macOS 13+ keeps a parallel "Background Items" database alongside
# classic LaunchAgents. Even when we wipe ~/Library/LaunchAgents/homebrew.mxcl.tailscale.plist
# in Step 5, the .btm file at
# ~/Library/Application Support/com.apple.backgroundtaskmanagementagent.backgrounditems.btm
# can still hold a Background-Items entry pointing at the (now-dead) plist
# path. Next login: macOS sees the entry, tries to bootstrap, finds the
# plist missing, and either silently regenerates it via brew's post-install
# hook (causing tailscaled respawn at next login) or refuses outright.
# `sfltool reset-login-items` is Apple's canonical API for wiping the
# entire .btm file. It DOES wipe all Login Items  the user re-adds what
# they want via System Settings  General  Login Items. Acceptable here
# because the goal of this script is "fresh baseline for Tailscale".
```

```
BTM="$HOME/Library/Application Support/com.apple.backgroundtaskmanagementagent/backgrounditems.btm"
if [ -f "$BTM" ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] found $BTM; would prompt sfltool reset-login-items"
  else
    confirm "Run 'sfltool reset-login-items' (wipes ALL Login Items; rebuild via System Settings  General  Login Items)?"
    if sfltool reset-login-items 2>&1; then
      echo "     Login Items drained"
    else
      echo "    ! sfltool reset-login-items failed (non-fatal; do manually: System Settings  General  Login Items)"
    fi
  fi
else
  echo "     no backgrounditems.btm present"
fi

# ---------- 6. brew install ----------
header "Step 6  brew install tailscale (fresh)"
if brew list tailscale >/dev/null 2>&1; then
  echo "     already installed  skip"
else
  confirm "Run 'brew install tailscale'?"
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew install tailscale"
  else
    brew install tailscale 2>&1
  fi
fi

# ---------- 6.5. Strip quarantine xattr from keg ----------
header "Step 6.5  Strip quarantine xattr from tailscale install"
# brew's freshly poured bottles don't carry com.apple.quarantine (the
# bottles are content-addressed and Homebrew-verified), but the upstream
# installer (Tailscale.app from App Store) and any manual download do.
# Strip the xattr from the entire keg as a defensive safety net so the
# freshly-installed tailscaled binary and any helper binaries are not
# Gatekeeper-blocked on first launch. Idempotent  no-op if no quarantine
# xattr is present.
quarantine_count="$(xattr -lr /opt/homebrew/Cellar/tailscale 2>/dev/null | grep -c '^[^ ]* com.apple.quarantine' || echo 0)"
if [ "${quarantine_count:-0}" -gt 0 ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] found $quarantine_count file(s) with com.apple.quarantine; would prompt xattr -dr"
  else
    confirm "Run 'xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale'? (strips quarantine xattr from $quarantine_count file(s); non-fatal if it errors)"
    if xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale 2>&1; then
      echo "     quarantine xattr stripped"
    else
      echo "    ! xattr -dr failed (non-fatal; open /opt/homebrew/Cellar/tailscale in Finder, right-click tailscaled  Open Anyway)"
    fi
  fi
else
  echo "     no com.apple.quarantine xattr on tailscale install"
fi

# ---------- 7. brew services start (NO sudo!) ----------
header "Step 7  Start the service as $USER (NO sudo, with multi-tier self-healing)"
if ! command -v tailscale >/dev/null 2>&1; then
  echo "     tailscale not on PATH after install  brew install errored; aborting"
  exit 4
fi

# Aliveness fan-out: ANY one of these is taken as proof of life. None of
# them alone is reliable  pgrep misses tailscaled if it does setproctitle
# or fork-detach; launchctl's reported PID can be a shim that exited
# within 1s; tailscale-cli hangs while macOS Local-Network permission is
# unresolved. Combined, they cover each other's blind spots.

tailscaled_api_up() {
  # Strategy A: tailscale CLI can talk to the local API. Canonical proof.
  # Bound at 8s  tailscale status itself returns quickly on success OR
  # failure, but if the post-install state is wedged it can hang.
  with_timeout 8 /opt/homebrew/bin/tailscale status >/dev/null 2>&1
}

tailscaled_proc_up() {
  # Strategy B: any process has 'tailscaled' substring in argv.
```

```
  pgrep -lf tailscaled >/dev/null 2>&1
}

tailscaled_lc_up() {
  # Strategy C: launchctl-managed service has a non-dash, non-zero PID.
  local lc_pid
  lc_pid="$(launchctl list 2>/dev/null | awk '$3 == "homebrew.mxcl.tailscale" {print $1; exit}')"
  if [ -n "$lc_pid" ] && [ "$lc_pid" != "-" ] && [ "$lc_pid" -gt 0 ] 2>/dev/null; then
    return 0
  fi
  return 1
}

tailscaled_is_alive() {
  tailscaled_api_up  && return 0
  tailscaled_proc_up && return 0
  tailscaled_lc_up   && return 0
  return 1
}

confirm "Run 'brew services start tailscale'?"
if [ "$DRY" = 1 ]; then
  echo "    [dry] brew services start tailscale"
  echo "    [dry] (then poll up to 30s for tailscaled_is_alive)"
  echo "    [dry] (recovery tier 1: launchctl kickstart)"
  echo "    [dry] (recovery tier 2: launchctl bootout + bootstrap of explicit plist)"
  echo "    [dry] (recovery tier 3: brew services restart (stop+start cycle))"
  echo "    [dry] (recovery tier 4: direct /opt/homebrew/bin/tailscaled spawn (bypasses launchd))"
else
  brew services start tailscale 2>&1
  # Primary poll  tailscaled_is_alive covers all 3 detection strategies.
  echo "     polling for tailscaled to come up (up to ~30s)"
  tailscaled_up=0
  for i in 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15; do
    if tailscaled_is_alive; then
      tailscaled_up=1
      echo "       tailscaled alive after ${i}*2s"
      break
    fi
    sleep 2
  done

  # Recovery tier 1: kickstart the launchd-managed service. This is the
  # cheapest and most-aligned-with-launchd fix when brew has loaded the
  # LaunchAgent but tailscaled didn't fork. Kickstart tells launchd
  # "re-spawn if not running" without forcing a teardown.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 1: launchctl kickstart"
    if launchctl kickstart -kp "gui/$UID/homebrew.mxcl.tailscale" 2>&1; then
      for i in 1 2 3 4 5 6 7 8 9 10; do
        if tailscaled_is_alive; then
          tailscaled_up=1
          echo "       tailscaled alive after kickstart (${i}*2s)"
          break
        fi
        sleep 2
      done
    else
      echo "      ! kickstart returned non-zero (LaunchAgent may not be loaded)"
    fi
  fi

  # Recovery tier 2: explicit plist bootout + bootstrap. Forces a fresh
  # launchd registration cycle for the brew-managed service, which fixes
  # the wider class of "brew says started but daemon never spawned"
  # wedge states that kickstart alone can't dig out of.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 2: launchctl bootout + bootstrap of explicit plist"
    launchctl bootout "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null || true
    sleep 1
    # Brew installs the canonical plist template at $HOMEBREW_PREFIX/Library/.
    # On Apple Silicon that's /opt/homebrew/Library/LaunchAgents/. On Intel
    # it's /usr/local/Library/LaunchAgents/. Detect which and pick the
    # right one so this script works across both architectures.
    if [ -f /opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
      plist_path="/opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
    elif [ -f /usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
      plist_path="/usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
    else
      plist_path=""
    fi
```

```
    if [ -n "$plist_path" ]; then
      if launchctl bootstrap "gui/$UID" "$plist_path" 2>&1; then
        for i in 1 2 3 4 5 6 7 8 9 10; do
          if tailscaled_is_alive; then
            tailscaled_up=1
            echo "       tailscaled alive after bootstrap (${i}*2s)"
            break
          fi
          sleep 2
        done
      else
        echo "      ! explicit bootstrap returned non-zero"
      fi
    else
      echo "      ! could not locate brew plist template; skipping bootstrap tier"
    fi
  fi

  # Recovery tier 3: brew services restart (forces a stop + start cycle,
  # which unsticks LaunchAgents with stale 'started' state from prior
  # half-cycles). This is the canonical brew-doc fix for 'services don't
  # actually start even though brew says started'.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 3: brew services restart (forces stop+start cycle)"
    if brew services restart tailscale 2>&1; then
      for i in 1 2 3 4 5 6 7 8 9 10; do
        if tailscaled_is_alive; then
          tailscaled_up=1
          echo "       tailscaled alive after brew restart (${i}*2s)"
          break
        fi
        sleep 2
      done
    else
      echo "      ! brew services restart returned non-zero"
    fi
  fi

  # Recovery tier 4: last-resort direct spawn of the daemon binary,
  # bypassing launchd entirely. If tailscaled crashes here too, we'll
  # capture its actual stderr in /tmp for the user to read (NOPASSWD
  # sudoers often have local-network permission denied  the captured
  # stderr makes it obvious what to fix in System Settings).
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 4: direct spawn (bypasses launchd; tail stderr to /tmp)"
    /opt/homebrew/bin/tailscaled >/tmp/nullexit-tailscaled-spawn.log 2>&1 &
    spawned_pid=$!
    disown "$spawned_pid" 2>/dev/null || true
    sleep 3
    if tailscaled_is_alive; then
      tailscaled_up=1
      echo "       tailscaled alive via direct spawn (pid=$spawned_pid)"
      echo "      (NOTE: not under launchd supervision  a reboot, re-run of"
      echo "       this script, or ./toggle.sh is required to recreate."
      echo "       Last 5 lines of stderr captured to /tmp/nullexit-tailscaled-spawn.log:)"
      tail -n 5 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/        /' || true
    fi
  fi

  if [ "$tailscaled_up" = 0 ]; then
    echo "     HARD FAIL: tailscaled refuses to stay alive across every recovery tier"
    echo "      launchctl-managed state (if available):"
    launchctl print "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null \
      | grep -E 'state =|last exit code =|runs =|path =' \
      | sed 's/^/        /' \
      || echo "        (no launchctl state available for homebrew.mxcl.tailscale)"
    echo "      Last 10 lines of /tmp/nullexit-tailscaled-spawn.log (if any):"
    tail -n 10 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/        /' || echo "        (no log)"
    echo "      Most likely remaining cause: macOS Local Network permission denied."
    echo "      System Settings  Privacy & Security  Local Network  toggle Tailscale ON."
    exit 7
  fi
fi
sleep 3

# ---------- 8. Verify ----------
header "Step 8  Verify fresh install"
echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{print $2, $3}' | tr '\n' ' ' | sed 's/ $//')"
echo "    tailscaled process:     $(pgrep -lf tailscaled 2>/dev/null | head -1 || echo 'NONE (check brew logs)')"
```

```
ts_ver=""
if [ -x /opt/homebrew/bin/tailscale ]; then
  ts_ver="$(/opt/homebrew/bin/tailscale version 2>/dev/null | head -1 | tr -d '\n')"
elif command -v tailscale >/dev/null 2>&1; then
  ts_ver="$(tailscale version 2>/dev/null | head -1 | tr -d '\n')"
fi
echo "    tailscale version:      ${ts_ver:-binary missing}"
ts_status=""
if [ -x /opt/homebrew/bin/tailscale ]; then
  ts_status="$(/opt/homebrew/bin/tailscale status 2>/dev/null | head -3)"
fi
if [ -n "$ts_status" ]; then
  echo "    tailscale status:"
  printf '%s\n' "$ts_status" | sed 's/^/      /'
else
  echo "    tailscale status:       (binary missing or errored)"
fi

if { [ "$DRY" = 0 ] && /opt/homebrew/bin/tailscale status 2>/dev/null | grep -q '^Tailscale is stopped'; }; then
  echo
  echo "     NOTICE: 'tailscale status' still reports stopped after a fresh install."
  echo "       Two equally likely causes:"
  echo "         (a) macOS hasn't yet shown the first-run 'Local Network' prompt for"
  echo "             Tailscale  open System Settings  Privacy & Security  Local"
  echo "             Network. Toggle Tailscale ON."
  echo "         (b) macOS denied Local Network permission at some prior point."
  echo "       Either way: System Settings  Privacy & Security  Local Network."
  echo "       If the entry isn't present, invoking any tailscale command will"
  echo "       re-trigger the prompt automatically."
fi

# ---------- 9. Done ----------
header "Done"
if [ "$DRY" = 1 ]; then
  echo "    (dry-mode finished; no follow-up commands were issued.)"
  exit 0
fi

cat <<'DONE'

   Manual follow-ups (all of these; the script stops short of running sudo
    tailscale up so you can verify the brew install is actually live first):

    1.  System Settings  Privacy & Security  Local Network  Tailscale ON
       macOS will normally pop the prompt the first time you run `tailscale`.
       If it didn't, trigger it via the menu bar  check.

    2.  Re-engage the tailnet + exit node:
           sudo tailscale up --ssh=true --accept-dns=false \
             --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
             --exit-node-allow-lan-access=true

    3.  Re-run: ./toggle.sh
       (re-installs the com.nullexit.wake-recovery LaunchAgent  wiped
       in Step 5 to clear the slate  and re-engages the gateway)

    4.  Final verification:
           curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
         expect: warp=on, ip=<Cloudflare IP>, colo=YYZ/YUL

    Container-side (gateway / 100.100.21.8) Tailscale is unaffected by this
    script. Other tailnet devices (S24, iPhone, Windows) are unaffected.

DONE
```

49 scripts/routing-fix.sh

```
#!/bin/sh
# routing-fix: Maintains routing for SOCKS5 proxy traffic through WARP (tun0)
#              and FORWARD rules for Tailscale exit-node return traffic.
#
# Runs inside the warp container's network namespace. This script ensures:
# 1. Table 200 has default via tun0 (for SOCKS5 proxy traffic from 172.18.0.2)
#    with the WARP endpoint exception via eth0 (prevents tunnel loop)
# 2. The CGNAT ip rule (100.64.0.0/10  table 52  tailscale routes) is maintained for Tailscale
```

```
# 3. FORWARD RELATED,ESTABLISHED rule allows return traffic for exit-node
#    forwarded connections (S24  tailscale0  eth0  host  internet)
#
# CRITICAL: Does NOT touch the main table's default route. Gluetun handles that
# internally via its own routing rules and WireGuard fwmark (0xca6c  table 51820).
# Overriding the main table default would create a loop: encrypted WireGuard
# packets exiting tun0 would match default dev tun0 and re-enter the tunnel.

set -e

trap 'echo "routing-fix: Caught signal, exiting..."; exit 0' TERM INT QUIT

sleep 15 &
wait $! || true

echo "routing-fix: Setting up routing for SOCKS5 proxy through WARP (tun0)..."

# Dynamically detect Docker subnet from eth0
DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')
DOCKER_GW=$(ip route show default | head -1 | awk '{print $3}')
WARP_ENDPOINT="${WARP_ENDPOINT}"
TAILSCALE_ALLOW_LAN_P2P="${TAILSCALE_ALLOW_LAN_P2P:-false}"
IP_BLOCKLIST_FILE="/userfilters/ip_blocklist.ipset"
LAST_IP_MTIME=""
echo "routing-fix: Docker network: ${DOCKER_NET}, gateway: ${DOCKER_GW}"

# Table 200: Used by ip rule 100 (from 172.18.0.2) for SOCKS5 proxy traffic
# WARP endpoint goes via eth0 to avoid tunnel loop, everything else via tun0
ip route add "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || \
  ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
ip route replace default dev tun0 table 200 2>/dev/null || true

# Main table: Just ensure the Docker subnet route exists via eth0
# (gluetun owns the default route here  do NOT touch it)
ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true

# FORWARD RELATED,ESTABLISHED: Allow return traffic for exit-node connections.
# The container has BOTH iptables (nftables backend) and iptables-legacy
# loaded. Gluetun's post-rules.txt uses the nftables backend, while Tailscale
# uses the legacy backend. Packets go through BOTH stacks, so we add the rule
# to both to ensure it survives regardless of which backend has policy DROP.
#
# Packet flow: S24 outgoing (tailscale0->eth0) ACCEPTed by first rule.
# Return traffic (eth0->eth0 via table 52 to host) needs RELATED,ESTABLISHED
# to match the conntrack entry created by the outgoing flow.
add_fwd_related_established() {
  # nftables backend (used by gluetun's post-rules.txt)
  if ! iptables -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
    iptables -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
  fi
  # legacy backend (used by Tailscale's ts-forward)
  if command -v iptables-legacy >/dev/null 2>&1; then
    if ! iptables-legacy -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
      iptables-legacy -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
    fi
  fi
}

add_fwd_related_established

# Policy Routing for Tailscale P2P:
# Tailscale sends its encrypted mesh UDP packets from port 41642.
# If these go out tun0 (WARP), Cloudflare's strict NAT drops the returning
# hole-punch packets, causing Tailscale to fail P2P and fall back to DERP relays (high latency).
# We mark packets originating from port 41642 and force them out the main table (eth0).
#
# LAN P2P is controlled by two sources (in priority order):
#   1. /app/.lan_p2p_detected  written by watcher.sh on macOS on every network change
#      using system_profiler SPAirPortDataType (802.1x detection) + AP isolation probe. Overrides .env.
#   2. TAILSCALE_ALLOW_LAN_P2P env var (from .env)  static fallback.
# When disabled, we explicitly DROP local RFC1918-destined Tailscale UDP so the
# phone cleanly falls back to DERP rather than getting poisoned by macOS gvproxy SNAT.
add_tailscale_p2p_bypass() {
  # Dynamic override from watcher.sh (network auto-detection) takes priority over .env
  local allow_p2p="${TAILSCALE_ALLOW_LAN_P2P:-false}"
  if [ -f "/app/.lan_p2p_detected" ]; then
    allow_p2p=$(cat "/app/.lan_p2p_detected" 2>/dev/null | tr -d '[:space:]')
  fi

  # Always allow P2P directly to the Docker host gateway (the Mac running this container).
  # The gateway container is always co-located with the host Mac on the same machine,
```

```
  # so direct Tailscale P2P via the Docker bridge is always safe and avoids DERP overhead.
  # This RETURN rule must be inserted BEFORE the general DROP rules so it takes priority.
  local default_gw
  default_gw=$(ip route show default | awk '{print $3}')
  if [ -n "$default_gw" ]; then
    if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null; then
      iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null || true
    fi
  fi

  # Also explicitly whitelist all physical IPs of the host Mac (e.g. 172.17.52.223).
  # The Tailscale control plane registers the Mac using its physical IP, so the
  # container will attempt P2P handshakes using that physical IP instead of the
  # Docker bridge IP. If this physical IP falls within 172.16.0.0/12, it would be dropped.
  if [ -f "/app/.host_ips" ]; then
    while IFS= read -r host_ip; do
      [ -z "$host_ip" ] && continue
      if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null; then
        iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null || true
      fi
    done < "/app/.host_ips"
  fi

  if [ "${allow_p2p}" != "true" ]; then
    # Drop Tailscale UDP packets destined for local subnets to prevent macOS gvproxy SNAT
    # endpoint poisoning. When disabled (campus/enterprise WPA2-Enterprise with AP isolation),
    # dropping these forces the S24 to use DERP relays reliably.
    for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
      if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; then
        iptables -t mangle -A OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null || true
      fi
    done
  else
    # Remove DROP rules if they exist (switching from restricted to allowed)
    for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
      while iptables -t mangle -D OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; do :; done
    done
  fi

  # Bypass STUN and P2P packets to public IPs (DERP relays and remote peers) out eth0
  if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null; then
    iptables -t mangle -A OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null || true
  fi
  if ! ip rule show | grep -q 'fwmark 0x8888'; then
    ip rule add fwmark 0x8888 lookup 254 pref 98 2>/dev/null || true
  fi
}

add_tailscale_p2p_bypass

add_onion_transparent_proxy() {
  for ipt in iptables iptables-legacy; do
    command -v "$ipt" >/dev/null 2>&1 || continue
    if ! $ipt -t nat -C PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null; then
      $ipt -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null || true
    fi
    if ! $ipt -C FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null; then
      $ipt -I FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null || true
    fi
  done
}

add_tailscale_p2p_bypass
add_onion_transparent_proxy

create_log_drop_chain() {
  local ipt="$1"
  local chain="$2"
  local prefix="$3"

  if ! $ipt -N "$chain" 2>/dev/null; then
    $ipt -F "$chain" 2>/dev/null || true
  fi
  $ipt -A "$chain" -j LOG --log-prefix "$prefix"
  $ipt -A "$chain" -j DROP 2>/dev/null || true
}

add_country_block() {
  # To add a country to the blocklist, simply define the 'BLOCKED_COUNTRIES' variable in .env with 2-letter ISO
      country codes.
  # You can find the country code by looking at the filename on http://www.ipdeny.com/ipblocks/data/countries/
```

```
  # For example, to block China (cn) and Russia (ru), set: BLOCKED_COUNTRIES="cn ru" in your .env file
  for zone in $BLOCKED_COUNTRIES; do
    set_name="block_${zone}"
    zone_upper=$(echo "$zone" | tr 'a-z' 'A-Z')
    chain_name="LOG_DROP_${zone_upper}"
    log_prefix="IP_BLOCK_${zone_upper}: "

    # Create the ipset if it doesn't exist
    if ! ipset list "$set_name" >/dev/null 2>&1; then
      ipset create "$set_name" hash:net 2>/dev/null || true
      echo "routing-fix: Downloading ${zone_upper} IPs for blocklist..."
      curl -sS "http://www.ipdeny.com/ipblocks/data/countries/${zone}.zone" | grep -v '^#' | while read -r subnet; do
        [ -n "$subnet" ] && ipset add "$set_name" "$subnet" 2>/dev/null || true
      done
    fi

    # Apply LOG_DROP rule to both backends
    for ipt in iptables iptables-legacy; do
      command -v "$ipt" >/dev/null 2>&1 || continue

      create_log_drop_chain "$ipt" "$chain_name" "$log_prefix"

      # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
      while $ipt -D FORWARD -m set --match-set "$set_name" dst -j DROP 2>/dev/null; do :; done

      if ! $ipt -C FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null; then
        $ipt -I FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null || true
      fi
    done
  done
}

add_ip_blocklist() {
  # Only reload when the compiled file has actually changed (mtime check).
  # This prevents a pointless ipset restore on every 30-second loop tick.
  if [ ! -f "$IP_BLOCKLIST_FILE" ]; then
    return 0
  fi

  CURRENT_MTIME=$(stat -c %Y "$IP_BLOCKLIST_FILE" 2>/dev/null || echo "0")
  if [ "$CURRENT_MTIME" != "$LAST_IP_MTIME" ]; then
    echo "routing-fix: IP blocklist changed, reloading..."
    # ipset restore handles the atomic swap internally:
    # create_new  populate  swap with live  destroy_new
    if ipset restore < "$IP_BLOCKLIST_FILE" 2>/dev/null; then
      LAST_IP_MTIME="$CURRENT_MTIME"
      echo "routing-fix: IP blocklist loaded ($(ipset list block_malicious | grep -c '^[0-9]') entries)."
    else
      echo "routing-fix: Warning: ipset restore failed. Retrying next cycle."
      return 1
    fi
  fi

  # Apply FORWARD rules in BOTH iptables backends.
  for ipt in iptables iptables-legacy; do
    command -v "$ipt" >/dev/null 2>&1 || continue

    create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_DST "IP_BLOCK_MALICIOUS_DST: "
    create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_SRC "IP_BLOCK_MALICIOUS_SRC: "

    # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
    while $ipt -D FORWARD -m set --match-set block_malicious dst -j DROP 2>/dev/null; do :; done
    while $ipt -D FORWARD -m set --match-set block_malicious src -j DROP 2>/dev/null; do :; done

    if ! $ipt -C FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null; then
      $ipt -I FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null || true
    fi

    if ! $ipt -C FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null; then
      $ipt -I FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null || true
    fi
  done
}

# Start background block logger (DNS + IP)
if ! pgrep -f "nullexit-logger" >/dev/null; then
  echo "routing-fix: Starting background block logger..."
  nullexit-logger &
fi

```

```
add_country_block
add_ip_blocklist

echo 'routing-fix: Routes applied.'

UNHEALTHY_COUNT=0

# Re-assert loop (every 30 seconds)
while true; do
  # Table 200 routes
  ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true

  # Tunnel health check
  if ! curl -m 3 -sS --interface tun0 https://cloudflare.com/cdn-cgi/trace >/dev/null 2>&1; then
    UNHEALTHY_COUNT=$((UNHEALTHY_COUNT + 1))
  else
    UNHEALTHY_COUNT=0
    if [ -f "/app/TUNNEL_FAILED_CLOSED.marker" ]; then
      rm -f "/app/TUNNEL_FAILED_CLOSED.marker"
    fi
  fi

  if [ "$UNHEALTHY_COUNT" -ge 3 ]; then
    echo "routing-fix: Tunnel health check failed $UNHEALTHY_COUNT times. Failing closed."
    ip route del default dev tun0 table 200 2>/dev/null || true
    ip route replace blackhole default table 200 2>/dev/null || true
    touch "/app/TUNNEL_FAILED_CLOSED.marker"
  else
    ip route replace default dev tun0 table 200 2>/dev/null || true
  fi

  # Main table: keep Docker subnet route
  ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true

  # CGNAT ip rule for Tailscale
  if ! ip rule show | grep -q '100.64.0.0/10'; then
    ip rule add to 100.64.0.0/10 lookup 52 pref 99 2>/dev/null || true
    echo 'routing-fix: CGNAT rule missing, re-injected'
  fi

  # FORWARD RELATED,ESTABLISHED: Allow return traffic in both iptables
  # backends. Gluetun may reset the nftables ruleset on VPN reconnect,
  # and Tailscale may reset the legacy ruleset on auth refresh.
  add_fwd_related_established

  # Ensure P2P packets bypass WARP to maintain low latency
  add_tailscale_p2p_bypass

  # Ensure transparent proxying for .onion is active
  add_onion_transparent_proxy

  # Enforce country blocklist
  add_country_block

  # Enforce IP blocklist (reloads automatically when file changes)
  add_ip_blocklist

  sleep 30 &
  wait $! || true
done
```

50 scripts/rule-compiler/go.mod

```
module nullexit/rule-compiler

go 1.22
```

51 scripts/rule-compiler/main.go

```
package main

import (
```

```
  "bufio"
  "bytes"
  "crypto/md5"
  "encoding/base64"
  "encoding/hex"
  "fmt"
  "io"
  "net/http"
  "net/netip"
  "os"
  "os/exec"
  "path/filepath"
  "regexp"
  "sort"
  "strings"
  "sync"
  "time"
)

const (
  BlackListPath = "black_list.txt"
  WhiteListPath = "white_list.txt"
  OutputPath    = "adguard/work/userfilters/compiled_rules.txt"
  IpOutputPath  = "adguard/work/userfilters/ip_blocklist.ipset"
  IpCacheDir    = "adguard/work/userfilters/cache/ip"
  CacheDir      = "adguard/work/userfilters/cache"
)

var CoreLists = []string{
  "https://raw.githubusercontent.com/anudeepND/blacklist/master/adservers.txt",
  "https://raw.githubusercontent.com/anudeepND/blacklist/master/facebook.txt",
  "https://raw.githubusercontent.com/crazy-max/WindowsSpyBlocker/master/data/hosts/spy.txt",
  "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.samsung.txt",
  "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.apple.txt",
  "https://abp.oisd.nl/basic/",
  "https://adaway.org/hosts.txt",
  "https://pgl.yoyo.org/adservers/serverlist.php?hostformat=adblockplus&showintro=1&mimetype=plaintext",
  "https://raw.githubusercontent.com/nextdns/cname-cloaking-blocklist/master/domains",
}

var MediumAdditions = []string{
  "https://raw.githubusercontent.com/lightswitch05/hosts/master/docs/lists/facebook-extended.txt",
  "https://someonewhocares.org/hosts/zero/hosts",
  "https://urlhaus.abuse.ch/downloads/hostfile/",
}

var HeavyAdditions = []string{
  "https://raw.githubusercontent.com/StevenBlack/hosts/master/hosts",
  "https://raw.githubusercontent.com/Perflyst/PiHoleBlocklist/master/SmartTV-AGH.txt",
  "https://raw.githubusercontent.com/DandelionSprout/adfilt/master/GameConsoleAdblockList.txt",
}

var IpBlockLists = []string{
  "https://feodotracker.abuse.ch/downloads/ipblocklist.txt",
  "https://www.spamhaus.org/drop/drop.txt",
  "https://rules.emergingthreats.net/fwrules/emerging-Block-IPs.txt",
  "https://cinsscore.com/list/ci-badguys.txt",
}

func getProfiles() map[string][]string {
  profiles := make(map[string][]string)

  light := append([]string(nil), CoreLists...)
  light = append(light, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/light.txt")
  profiles["light"] = light

  medium := append([]string(nil), CoreLists...)
  medium = append(medium, MediumAdditions...)
  medium = append(medium, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/multi.txt")
  profiles["medium"] = medium

  heavy := append([]string(nil), CoreLists...)
  heavy = append(heavy, MediumAdditions...)
  heavy = append(heavy, HeavyAdditions...)
  heavy = append(heavy, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/pro.txt")
  profiles["heavy"] = heavy

  return profiles
}

func loadEnvProfile() string {
```

```
  profile := "medium"

  if bytesData, err := os.ReadFile(".env"); err == nil {
    scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
    for scanner.Scan() {
      line := strings.TrimSpace(scanner.Text())
      if line != "" && !strings.HasPrefix(line, "#") {
        parts := strings.SplitN(line, "=", 2)
        if len(parts) == 2 {
          key := strings.TrimSpace(parts[0])
          val := strings.TrimSpace(parts[1])
          if key == "GATEWAY_RULE_PROFILE" {
            val = strings.Trim(val, "\"'")
            profile = strings.ToLower(val)
          }
        }
      }
    }
  } else if !os.IsNotExist(err) {
    fmt.Printf("Warning: Failed to read .env file for profile configuration (%v). Using default.\n", err)
  }

  if envVal := os.Getenv("GATEWAY_RULE_PROFILE"); envVal != "" {
    profile = strings.ToLower(envVal)
  }

  profiles := getProfiles()
  if _, exists := profiles[profile]; !exists {
    fmt.Printf("Warning: Profile '%s' is invalid. Falling back to 'medium'.\n", profile)
    profile = "medium"
  }
  return profile
}

func loadDomains(filepath string) map[string]bool {
  domains := make(map[string]bool)
  if bytesData, err := os.ReadFile(filepath); err == nil {
    scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
    for scanner.Scan() {
      line := strings.TrimSpace(scanner.Text())
      if line != "" && !strings.HasPrefix(line, "#") {
        domains[strings.ToLower(line)] = true
      }
    }
  }
  return domains
}

func parseDomainsFromContent(content string) map[string]bool {
  domains := make(map[string]bool)
  scanner := bufio.NewScanner(strings.NewReader(content))
  ipRegex := regexp.MustCompile(`^\d{1,3}\.\d{1,3}\.\d{1,3}\.\d{1,3}$`)
  agRegex := regexp.MustCompile(`^\|\|([a-z0-9.-]+)\^$`)
  domainRegex := regexp.MustCompile(`^[a-z0-9.-]+$`)

  for scanner.Scan() {
    line := strings.TrimSpace(scanner.Text())
    if line == "" || strings.HasPrefix(line, "!") || strings.HasPrefix(line, "[") {
      continue
    }
    parts := strings.SplitN(line, "#", 2)
    line = strings.TrimSpace(parts[0])
    if line == "" {
      continue
    }

    fields := strings.Fields(line)
    var rule string
    if len(fields) >= 2 {
      if ipRegex.MatchString(fields[0]) {
        rule = fields[1]
      } else {
        rule = fields[0]
      }
    } else {
      rule = fields[0]
    }

    rule = strings.TrimSpace(strings.ToLower(rule))

    var domain string
```

```
    if m := agRegex.FindStringSubmatch(rule); m != nil {
      domain = m[1]
    } else if domainRegex.MatchString(rule) {
      domain = rule
    } else {
      domain = rule
    }

    if domain != "" && domain != "localhost" && domain != "0.0.0.0" && domain != "127.0.0.1" && domain != "
     broadcasthost" {
      domains[domain] = true
    }
  }
  return domains
}

// WARNING: Parsing AdGuardHome.yaml using regex instead of a proper YAML parser
// is brittle and assumes the exact format currently emitted by AdGuard.
// If an AdGuard version bump updates the configuration schema format,
// this parser will fail silently or loudly and should be refactored to use gopkg.in/yaml.v3.
func getAdguardNativeLists() []string {
  yamlPath := "adguard/conf/AdGuardHome.yaml"
  var enabledUrls []string
  if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
    return enabledUrls
  }

  contentBytes, err := os.ReadFile(yamlPath)
  if err != nil {
    fmt.Printf("Warning: Failed to parse AdGuardHome.yaml for native lists (%v)\n", err)
    return enabledUrls
  }
  content := string(contentBytes)

  filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*?)(?:^[a-zA-Z_]+:|\z)`)
  match := filtersRegex.FindStringSubmatch(content)
  if len(match) > 1 {
    filtersText := match[1]
    blocks := strings.Split(filtersText, "\n  - ")
    urlRegex := regexp.MustCompile(`url:\s*(\S+)`)

    for _, block := range blocks {
      if strings.TrimSpace(block) == "" {
        continue
      }
      if strings.Contains(block, "enabled: true") {
        urlMatch := urlRegex.FindStringSubmatch(block)
        if len(urlMatch) > 1 {
          url := urlMatch[1]
          if !strings.Contains(url, "compiled_rules.txt") {
            enabledUrls = append(enabledUrls, url)
          }
        }
      }
    }
  }
  return enabledUrls
}

func getAdguardCompiledRulesCachePath() string {
  yamlPath := "adguard/conf/AdGuardHome.yaml"
  if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
    return ""
  }

  contentBytes, err := os.ReadFile(yamlPath)
  if err != nil {
    fmt.Printf("Warning: Failed to resolve compiled_rules.txt cache path (%v)\n", err)
    return ""
  }
  content := string(contentBytes)

  filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*?)(?:^[a-zA-Z_]+:|\z)`)
  match := filtersRegex.FindStringSubmatch(content)
  if len(match) > 1 {
    blocks := strings.Split(match[1], "\n  - ")
    idRegex := regexp.MustCompile(`id:\s*(\d+)`)
    for _, block := range blocks {
      if !strings.Contains(block, "compiled_rules.txt") {
        continue
      }
```

```
      idMatch := idRegex.FindStringSubmatch(block)
      if len(idMatch) > 1 {
        return fmt.Sprintf("adguard/work/data/filters/%s.txt", idMatch[1])
      }
    }
  }
  return ""
}

func handleFetchError(urlStr string, cacheFile string, err error) map[string]bool {
  if _, statErr := os.Stat(cacheFile); statErr == nil {
    fmt.Printf(" -> Warning: Failed to fetch %s (%v). Falling back to expired local cache.\n", urlStr, err)
    content, readErr := os.ReadFile(cacheFile)
    if readErr == nil {
      domains := parseDomainsFromContent(string(content))
      fmt.Printf(" -> Loaded from expired cache (%d domains).\n", len(domains))
      return domains
    }
    fmt.Printf(" -> Warning: Failed to read expired cache (%v). Using local lists only.\n", readErr)
  } else {
    fmt.Printf(" -> Warning: Failed to fetch %s (%v). Using local lists only for this source.\n", urlStr, err)
  }
  return make(map[string]bool)
}

func fetchRemoteDomains(urlStr string) map[string]bool {
  os.MkdirAll(CacheDir, 0755)

  hash := md5.Sum([]byte(urlStr))
  urlHash := hex.EncodeToString(hash[:])
  cacheFile := filepath.Join(CacheDir, fmt.Sprintf("%s.txt", urlHash))

  if stat, err := os.Stat(cacheFile); err == nil {
    fileAge := time.Since(stat.ModTime()).Seconds()
    if fileAge < 86400 {
      fmt.Printf("Loading remote blacklist from cache: %s\n", urlStr)
      content, err := os.ReadFile(cacheFile)
      if err == nil {
        domains := parseDomainsFromContent(string(content))
        fmt.Printf(" -> Loaded from cache (copy is %.1f hours old, %d domains).\n", fileAge/3600.0, len(domains
    ))
        return domains
      }
      fmt.Printf(" -> Warning: Failed to read cache file %s (%v). Re-fetching...\n", cacheFile, err)
    }
  }

  fmt.Printf("Fetching remote blacklist from: %s ...\n", urlStr)

  req, err := http.NewRequest("GET", urlStr, nil)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }
  req.Header.Set("User-Agent", "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7)")

  client := &http.Client{Timeout: 15 * time.Second}
  resp, err := client.Do(req)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }
  defer resp.Body.Close()

  if resp.StatusCode != 200 {
    return handleFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
  }

  bodyBytes, err := io.ReadAll(resp.Body)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }

  content := string(bodyBytes)
  domains := parseDomainsFromContent(content)

  if len(domains) < 10 {
    return handleFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only found %d domains. Possible
      404 or bad URL", len(domains)))
  }

  err = os.WriteFile(cacheFile, bodyBytes, 0644)
  if err != nil {
```

```
    fmt.Printf(" -> Warning: Failed to save cache file (%v)\n", err)
  }

  fmt.Printf(" -> Successfully fetched and cached %d domains.\n", len(domains))
  return domains
}

func optimizeSubdomains(domains map[string]bool, listName string) map[string]bool {
  fmt.Printf("Optimizing %s by removing redundant subdomains...\n", listName)
  rawCount := len(domains)
  if rawCount == 0 {
    return make(map[string]bool)
  }

  optimized := make(map[string]bool)
  for domain := range domains {
    if strings.ContainsAny(domain, "|^@$/") {
      optimized[domain] = true
      continue
    }

    parts := strings.Split(domain, ".")
    hasParent := false

    for i := 1; i < len(parts)-1; i++ {
      parent := strings.Join(parts[i:], ".")
      if domains[parent] {
        hasParent = true
        break
      }
    }

    if !hasParent {
      optimized[domain] = true
    }
  }

  saved := rawCount - len(optimized)
  reduction := float64(0)
  if rawCount > 0 {
    reduction = (float64(saved) / float64(rawCount)) * 100
  }
  fmt.Printf(" -> Reduced %s from %d to %d domains (-%d / %.1f%% reduction).\n", listName, rawCount, len(
     optimized), saved, reduction)
  return optimized
}

func parseIpsFromContent(content string) map[string]bool {
  ips := make(map[string]bool)
  scanner := bufio.NewScanner(strings.NewReader(content))

  for scanner.Scan() {
    line := strings.TrimSpace(scanner.Text())
    if line == "" || strings.HasPrefix(line, "#") || strings.HasPrefix(line, ";") {
      continue
    }

    fields := strings.Fields(line)
    if len(fields) > 0 {
      token := strings.TrimRight(fields[0], ";,")
      if _, err := netip.ParsePrefix(token); err == nil {
        ips[token] = true
      } else if _, err := netip.ParseAddr(token); err == nil {
        ips[token] = true
      }
    }
  }
  return ips
}

func handleIpFetchError(urlStr string, cacheFile string, err error) map[string]bool {
  if _, statErr := os.Stat(cacheFile); statErr == nil {
    fmt.Printf(" -> Warning: fetch failed (%v). Falling back to stale cache.\n", err)
    content, readErr := os.ReadFile(cacheFile)
    if readErr == nil {
      ips := parseIpsFromContent(string(content))
      fmt.Printf(" -> %d IPs loaded from stale cache.\n", len(ips))
      return ips
    }
    fmt.Printf(" -> Warning: stale cache also failed (%v).\n", readErr)
  } else {
```

```
    fmt.Printf(" -> Warning: %s failed (%v). No cache available. Skipping.\n", urlStr, err)
  }
  return make(map[string]bool)
}

func fetchRemoteIps(urlStr string) map[string]bool {
  os.MkdirAll(IpCacheDir, 0755)

  hash := md5.Sum([]byte(urlStr))
  urlHash := hex.EncodeToString(hash[:])
  cacheFile := filepath.Join(IpCacheDir, fmt.Sprintf("%s.txt", urlHash))

  if stat, err := os.Stat(cacheFile); err == nil {
    fileAge := time.Since(stat.ModTime()).Seconds()
    if fileAge < 86400 {
      fmt.Printf("Loading IP feed from cache: %s\n", urlStr)
      content, err := os.ReadFile(cacheFile)
      if err == nil {
        ips := parseIpsFromContent(string(content))
        fmt.Printf(" -> %d IPs (cache is %.1fh old).\n", len(ips), fileAge/3600.0)
        return ips
      }
      fmt.Printf(" -> Warning: cache read failed (%v). Re-fetching...\n", err)
    }
  }

  fmt.Printf("Fetching IP feed: %s ...\n", urlStr)
  req, err := http.NewRequest("GET", urlStr, nil)
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }
  req.Header.Set("User-Agent", "Mozilla/5.0")

  client := &http.Client{Timeout: 15 * time.Second}
  resp, err := client.Do(req)
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }
  defer resp.Body.Close()

  if resp.StatusCode != 200 {
    return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
  }

  bodyBytes, err := io.ReadAll(resp.Body)
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }

  content := string(bodyBytes)
  ips := parseIpsFromContent(content)
  if len(ips) < 1 {
    return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only %d IPs found. Possible
     404", len(ips)))
  }

  err = os.WriteFile(cacheFile, bodyBytes, 0644)
  if err != nil {
    fmt.Printf(" -> Warning: cache write failed (%v)\n", err)
  }

  fmt.Printf(" -> %d IPs fetched and cached.\n", len(ips))
  return ips
}

func compileIpBlocklist() int {
  fmt.Println("\n IP Blocklist Compilation ")

  allIps := make(map[string]bool)
  var mu sync.Mutex
  var wg sync.WaitGroup

  maxThreads := 8
  if len(IpBlockLists) < 8 {
    maxThreads = len(IpBlockLists)
  }
  semaphore := make(chan struct{}, maxThreads)

  for _, urlStr := range IpBlockLists {
    wg.Add(1)
    go func(u string) {
```

```
      defer wg.Done()
      semaphore <- struct{}{}
      defer func() { <-semaphore }()

      ips := fetchRemoteIps(u)
      mu.Lock()
      for ip := range ips {
        allIps[ip] = true
      }
      mu.Unlock()
    }(urlStr)
  }
  wg.Wait()

  fmt.Printf("Total unique entries before filtering: %d\n", len(allIps))

  reservedStr := []string{
    "10.0.0.0/8",
    "172.16.0.0/12",
    "192.168.0.0/16",
    "127.0.0.0/8",
    "169.254.0.0/16",
    "100.64.0.0/10",
    "0.0.0.0/8",
  }
  var reserved []netip.Prefix
  for _, s := range reservedStr {
    p, _ := netip.ParsePrefix(s)
    reserved = append(reserved, p)
  }

  cleanIps := make(map[string]bool)
  for entry := range allIps {
    var prefix netip.Prefix
    p, err := netip.ParsePrefix(entry)
    if err != nil {
      addr, err2 := netip.ParseAddr(entry)
      if err2 != nil {
        continue
      }
      prefix = netip.PrefixFrom(addr, addr.BitLen())
    } else {
      prefix = p
    }

    overlaps := false
    for _, r := range reserved {
      if r.Overlaps(prefix) {
        overlaps = true
        break
      }
    }
    if !overlaps {
      cleanIps[entry] = true
    }
  }

  removed := len(allIps) - len(cleanIps)
  if removed > 0 {
    fmt.Printf(" -> Removed %d private/reserved entries.\n", removed)
  }
  fmt.Printf(" -> Final IP blocklist: %d entries.\n", len(cleanIps))

  os.MkdirAll(filepath.Dir(IpOutputPath), 0755)

  f, err := os.Create(IpOutputPath)
  if err != nil {
    fmt.Printf("Error writing IP output file: %v\n", err)
    return len(cleanIps)
  }
  defer f.Close()

  f.WriteString("# nullexit Compiled IP Blocklist\n")
  f.WriteString("# Sources: Feodo Tracker, Spamhaus DROP/EDROP, Emerging Threats, CINS\n")
  f.WriteString(fmt.Sprintf("# Entries: %d\n\n", len(cleanIps)))

  f.WriteString("create block_malicious_new hash:net maxelem 200000 -exist\n")

  var sortedIps []string
  for ip := range cleanIps {
    sortedIps = append(sortedIps, ip)
```

```
  }
  sort.Strings(sortedIps)

  for _, ip := range sortedIps {
    f.WriteString(fmt.Sprintf("add block_malicious_new %s -exist\n", ip))
  }

  f.WriteString("create block_malicious hash:net maxelem 200000 -exist\n")
  f.WriteString("swap block_malicious block_malicious_new\n")
  f.WriteString("destroy block_malicious_new\n")

  fmt.Printf("IP blocklist written to %s\n", IpOutputPath)
  return len(cleanIps)
}

func main() {
  profile := loadEnvProfile()
  fmt.Printf("Active memory profile: '%s'\n", strings.ToUpper(profile))

  localBlackList := loadDomains(BlackListPath)
  whiteList := loadDomains(WhiteListPath)

  blackList := make(map[string]bool)
  for k := range localBlackList {
    blackList[k] = true
  }

  profiles := getProfiles()
  urlsToFetch := profiles[profile]

  fmt.Printf("\nStarting concurrent downloads for %d lists...\n", len(urlsToFetch))
  startTime := time.Now()

  var mu sync.Mutex
  var wg sync.WaitGroup

  maxThreads := 16
  if len(urlsToFetch) < 16 {
    maxThreads = len(urlsToFetch)
  }
  semaphore := make(chan struct{}, maxThreads)

  for _, urlStr := range urlsToFetch {
    wg.Add(1)
    go func(u string) {
      defer wg.Done()
      semaphore <- struct{}{}
      defer func() { <-semaphore }()

      domains := fetchRemoteDomains(u)
      mu.Lock()
      for d := range domains {
        blackList[d] = true
      }
      mu.Unlock()
    }(urlStr)
  }
  wg.Wait()

  fmt.Printf("Finished concurrent downloads in %.2f seconds using %d threads.\n", time.Since(startTime).Seconds
     (), maxThreads)

  adguardNativeUrls := getAdguardNativeLists()
  var adguardNativeDomains map[string]bool
  if len(adguardNativeUrls) > 0 {
    fmt.Printf("\nFetching %d AdGuard native list(s) for deduplication...\n", len(adguardNativeUrls))
    adguardNativeDomains = make(map[string]bool)

    maxNativeThreads := 4
    if len(adguardNativeUrls) < 4 {
      maxNativeThreads = len(adguardNativeUrls)
    }
    nativeSemaphore := make(chan struct{}, maxNativeThreads)

    for _, urlStr := range adguardNativeUrls {
      wg.Add(1)
      go func(u string) {
        defer wg.Done()
        nativeSemaphore <- struct{}{}
        defer func() { <-nativeSemaphore }()

```

```
        domains := fetchRemoteDomains(u)
        mu.Lock()
        for d := range domains {
          adguardNativeDomains[d] = true
        }
        mu.Unlock()
      }(urlStr)
    }
    wg.Wait()

    if len(adguardNativeDomains) > 0 {
      fmt.Println("Deduplicating compiled rules against AdGuard native lists...")
      originalSize := len(blackList)
      for d := range adguardNativeDomains {
        delete(blackList, d)
      }
      fmt.Printf(" -> Removed %d redundant rules already covered by AdGuard.\n", originalSize-len(blackList))
    }
  }

  var contradictions []string
  for d := range whiteList {
    if blackList[d] {
      contradictions = append(contradictions, d)
    }
  }

  if len(contradictions) > 0 {
    fmt.Printf("Detected %d contradictions. Whitelist taking priority.\n", len(contradictions))
    sort.Strings(contradictions)
    for _, domain := range contradictions {
      if localBlackList[domain] {
        fmt.Printf(" -> Removed local blacklist domain: %s\n", domain)
      }
      delete(blackList, domain)
    }
  }

  blackList = optimizeSubdomains(blackList, "blacklist")
  whiteList = optimizeSubdomains(whiteList, "whitelist")

  os.MkdirAll(filepath.Dir(OutputPath), 0755)
  outFile, err := os.Create(OutputPath)
  if err != nil {
    fmt.Printf("Error creating output file: %v\n", err)
    return
  }
  defer outFile.Close()

  outFile.WriteString("! Custom Compiled Rules (Auto-Generated)\n")
  outFile.WriteString(fmt.Sprintf("! Memory Profile: %s\n", strings.ToUpper(profile)))
  outFile.WriteString(fmt.Sprintf("! Total Block Rules: %d\n", len(blackList)))
  outFile.WriteString(fmt.Sprintf("! Native AdGuard Rules: %d\n", len(adguardNativeDomains)))
  outFile.WriteString(fmt.Sprintf("! Total Whitelist Rules: %d\n\n", len(whiteList)))

  outFile.WriteString("! --- Blacklist Rules ---\n")

  var sortedBlacklist []string
  for d := range blackList {
    sortedBlacklist = append(sortedBlacklist, d)
  }
  sort.Strings(sortedBlacklist)

  for _, domain := range sortedBlacklist {
    if strings.Contains(domain, "$") {
      parts := strings.SplitN(domain, "$", 2)
      dom, mod := parts[0], parts[1]
      if strings.HasPrefix(dom, "||") {
        if strings.HasSuffix(dom, "^") {
          outFile.WriteString(fmt.Sprintf("%s$%s\n", dom, mod))
        } else {
          outFile.WriteString(fmt.Sprintf("%s^$%s\n", dom, mod))
        }
      } else {
        outFile.WriteString(fmt.Sprintf("||%s^$%s\n", dom, mod))
      }
    } else if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|"
    ) || strings.HasSuffix(domain, "^") {
      outFile.WriteString(fmt.Sprintf("%s\n", domain))
    } else {
      outFile.WriteString(fmt.Sprintf("||%s^\n", domain))
```

```
    }
  }

  outFile.WriteString("\n! --- Whitelist Rules ---\n")
  var sortedWhitelist []string
  for d := range whiteList {
    sortedWhitelist = append(sortedWhitelist, d)
  }
  sort.Strings(sortedWhitelist)

  for _, domain := range sortedWhitelist {
    if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|") ||
     strings.HasSuffix(domain, "^") || strings.HasPrefix(domain, "@@") {
      rule := domain
      if !strings.HasPrefix(domain, "@@") {
        rule = "@@" + domain
      }
      outFile.WriteString(fmt.Sprintf("%s\n", rule))
    } else {
      outFile.WriteString(fmt.Sprintf("@@||%s^\n", domain))
    }
  }

  fmt.Printf("\nSuccessfully compiled %d block rules and %d allow rules to %s\n", len(blackList), len(whiteList
     ), OutputPath)

  cachedFilterPath := getAdguardCompiledRulesCachePath()
  if cachedFilterPath != "" {
    input, err := os.ReadFile(OutputPath)
    if err == nil {
      err = os.WriteFile(cachedFilterPath, input, 0644)
      if err == nil {
        fmt.Printf("Updated AdGuard filter cache: %s\n", cachedFilterPath)
      } else {
        fmt.Printf("Warning: Could not update AdGuard filter cache %s (%v)\n", cachedFilterPath, err)
      }
    } else {
      fmt.Printf("Warning: Could not read OutputPath to copy to cache (%v)\n", err)
    }
  } else {
    fmt.Println("Warning: Could not determine compiled_rules.txt cache path from AdGuardHome.yaml; skipping
     cache update.")
  }

  adguardReachable := false

  dockerComposeYamlExists := false
  if _, err := os.Stat("docker-compose.yml"); err == nil {
    dockerComposeYamlExists = true
  }
  dockerPath, err := exec.LookPath("docker")
  if err == nil && dockerComposeYamlExists {
    cmd := exec.Command(dockerPath, "compose", "ps", "--status", "running", "-q", "adguardhome")
    out, err := cmd.Output()
    if err == nil && strings.TrimSpace(string(out)) != "" {
      fmt.Println("Force-recreating AdGuard Home container to drop persisted filter state...")
      upCmd := exec.Command(dockerPath, "compose", "up", "-d", "--force-recreate", "--no-deps", "adguardhome")
      err := upCmd.Run()
      if err == nil {
        fmt.Println("AdGuard Home force-recreated successfully.")
        time.Sleep(8 * time.Second)
        adguardReachable = true
      } else {
        fmt.Println("Failed to restart AdGuard Home. Is it running?")
      }
    }
  }

  if adguardReachable {
    agPass := "nullexit"
    if bytesData, err := os.ReadFile(".env"); err == nil {
      scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
      for scanner.Scan() {
        line := scanner.Text()
        if strings.HasPrefix(line, "ADGUARD_PASSWORD=") {
          parts := strings.SplitN(line, "=", 2)
          if len(parts) == 2 {
            agPass = strings.Trim(strings.TrimSpace(parts[1]), "\"'")
          }
          break
        }
```

```
      }
    }

    creds := base64.StdEncoding.EncodeToString([]byte(fmt.Sprintf("admin:%s", agPass)))
    req, err := http.NewRequest("POST", "http://127.0.0.1:3000/control/filtering/refresh", bytes.NewBuffer([]byte("{}")))
    if err == nil {
      req.Header.Set("Authorization", fmt.Sprintf("Basic %s", creds))
      req.Header.Set("Content-Type", "application/json")
      client := &http.Client{Timeout: 15 * time.Second}
      resp, err := client.Do(req)
      if err == nil {
        fmt.Printf("Triggered AdGuard filter refresh via REST API (HTTP=%d).\n", resp.StatusCode)
        resp.Body.Close()
      } else {
        fmt.Printf("Warning: AdGuard filter refresh API call failed (%v). New whitelist will not load until next refresh tick (~24h) or manual refresh.\n", err)
      }
    }
  }

  compileIpBlocklist()
}
```

52 scripts/setup-common.sh

```
#!/bin/bash
#  1. Docker
step "Checking Docker"

if ! command -v docker >> output.log 2>&1; then
    echo ""
    if [[ "$OSTYPE" == "darwin"* ]]; then
        echo "  Docker is not installed."
        # Note: Homebrew aggressively rolls forward and discourages version pinning.
        # This setup is confirmed stable on Colima v0.10.3 and Docker v29.6.1.
        read -rp "  Would you like to automatically install Colima & Docker via Homebrew? [y/N]: " INSTALL_DOCKER
        if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
            step "Installing Colima and Docker..."
            brew install colima docker docker-compose
            step "Starting Colima VM..."
            colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged
        else
            echo ""
            echo "  Please install Docker manually. Two options:"
            echo "  A) Colima (Recommended): brew install colima docker docker-compose && colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged"
            echo "  B) Docker Desktop: https://www.docker.com/products/docker-desktop/"
            die "Install Docker, then re-run this script."
        fi
    else
        echo "  Docker is not installed."
        # Note: The Docker convenience script always fetches the latest stable release.
        # This setup is confirmed stable on Docker Engine v29.6.1.
        read -rp "  Would you like to automatically install Docker? (Requires sudo) [y/N]: " INSTALL_DOCKER
        if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
            step "Installing Docker..."
            curl -fsSL https://get.docker.com | sh
            sudo usermod -aG docker "$USER"
            echo -e "\n\033[0;32m[OK] Docker installed successfully!\033[0m"
            echo -e "\033[0;33m[!] CRITICAL: You must log out of your SSH session and log back in for Docker permissions to apply.\033[0m"
            die "Please log out, log back in, and re-run setup.sh."
        else
            echo "  Please install Docker manually:"
            echo "       curl -fsSL https://get.docker.com | sh"
            echo "       sudo usermod -aG docker \$USER && newgrp docker"
            die "Install Docker, then re-run this script."
        fi
    fi
fi

if ! docker info >> output.log 2>&1; then
    echo ""
    if [[ "$OSTYPE" == "darwin"* ]]; then
```

```
        echo "  Docker is installed but not running."
        echo "  If using Colima:        colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged"
        echo "  If using Docker Desktop: open the app."
    else
        echo "  Docker is installed but not running."
        echo "  Run: sudo systemctl start docker"
    fi
    die "Start Docker, then re-run this script."
fi

if ! docker compose version >> output.log 2>&1; then
    die "Docker Compose v2 not found. Update Docker or install the compose plugin:\n  https://docs.docker.com/compose/install/"
fi

ok "Docker $(docker --version | grep -oE '[0-9.]+' | head -1) is running."

#  1.5 Python 3
step "Checking Python 3"

if ! command -v python3 >> output.log 2>&1; then
    echo ""
    if [[ "$OSTYPE" == "darwin"* ]]; then
        echo "  Python 3 is not installed."
        echo "  macOS includes it via Xcode Command Line Tools, or you can install it via Homebrew:"
        echo "       brew install python3"
        die "Install Python 3, then re-run this script."
    else
        echo "  Python 3 is not installed."
        echo "  Please install it using your system's package manager. For example:"
        echo "       Debian/Ubuntu:  sudo apt install python3"
        echo "       Fedora:         sudo dnf install python3"
        echo "       Arch:           sudo pacman -S python"
        die "Install Python 3, then re-run this script."
    fi
fi

# Note: This setup is confirmed stable on Python 3.9.6 (macOS default).
PYTHON_VER=$(python3 --version 2>&1 | head -n 1)
ok "$PYTHON_VER is installed."

#  2. Tailscale (host)
step "Checking Tailscale (host)"
# The host needs its own Tailscale installation  separate from the Docker
# container  so Toggle-Gateway can run 'tailscale down' before stopping
# containers and 'tailscale up' after starting them. Without this, the
# exit-node deadlock that setup.sh is designed to prevent can still occur.

# Resolve the tailscale binary: CLI (brew/package) or bundled macOS app
TAILSCALE_BIN=""
if command -v tailscale >> output.log 2>&1; then
    TAILSCALE_BIN="tailscale"
elif [[ -x "/Applications/Tailscale.app/Contents/MacOS/Tailscale" ]]; then
    TAILSCALE_BIN="/Applications/Tailscale.app/Contents/MacOS/Tailscale"
fi

RECOMMENDED_TS_VERSION="1.98.5"

if [[ -z "$TAILSCALE_BIN" ]]; then
    warn "Tailscale not found on host  installing..."

    if [[ "$OSTYPE" == "darwin"* ]]; then
        if command -v brew >> output.log 2>&1; then
            step "Installing Tailscale CLI formula via Homebrew (standalone, no .app GUI)..."
            # Note: We install Tailscale via Homebrew, targeting the stable version (recommended: 1.98.5)
            brew install tailscale -q
            TAILSCALE_BIN="tailscale"

            step "Starting tailscaled as a per-user LaunchAgent (auto-starts at login)..."
            if brew services start tailscale 2>&1 | grep -qE 'Successfully|started'; then
                ok "tailscaled LaunchAgent registered."
            else
                warn "Could not auto-start tailscaled. Run 'brew services start tailscale' manually."
                warn "For a system-wide LaunchDaemon that runs before login, run instead:"
                warn "    sudo brew services start tailscale"
            fi
            echo ""
        else
            die "Homebrew not found. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION}):\n  https://tailscale.com/download/mac\n  Then re-run this script."
```

```
        fi

    elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
        curl -fsSL https://tailscale.com/install.sh | sh
        sudo systemctl enable --now tailscaled 2>> output.log || true
        TAILSCALE_BIN="tailscale"

    else
        die "Unsupported OS. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION}):\n
     https://tailscale.com/download\n  Then re-run this script."
    fi

    TAILSCALE_VER=$(${TAILSCALE_BIN} version 2>> output.log | head -n 1)
    if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
        warn "Tailscale installed, but version ($TAILSCALE_VER) differs from recommended version (
     $RECOMMENDED_TS_VERSION)."
    else
        ok "Tailscale installed (version $TAILSCALE_VER)."
    fi
else
    TAILSCALE_VER=$(${TAILSCALE_BIN} version 2>> output.log | head -n 1)
    if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
        warn "Tailscale is already installed ($TAILSCALE_VER), but recommended version is
     $RECOMMENDED_TS_VERSION for host/container consistency."
    else
        ok "Tailscale is already installed at recommended version ($TAILSCALE_VER)."
    fi
fi

# Authenticate the host machine if it isn't already connected.
# This is an interactive step  it opens a browser login URL.
# We do it now, before containers start, so Toggle-Gateway can
# safely run 'tailscale down/up' from day one.
if ! ${TAILSCALE_BIN} status >> output.log 2>&1; then
    echo ""
    echo "  Your host machine needs to join your Tailscale network."
    echo "  A browser window or login URL will appear  authenticate, then"
    echo "  return here. The script will continue automatically."
    echo ""
    if [[ "$OSTYPE" == "linux-gnu"* ]]; then
        sudo ${TAILSCALE_BIN} up
    else
        ${TAILSCALE_BIN} up
    fi
    ok "Host machine connected to Tailscale."
else
    ok "Host machine is already connected to Tailscale."
fi

#  3. wgcf
step "Checking wgcf (Cloudflare WARP key tool)"

if ! command -v wgcf >> output.log 2>&1; then
    warn "wgcf not found  installing..."

    if [[ "$OSTYPE" == "darwin"* ]]; then
        command -v brew >> output.log 2>&1 || die "Homebrew not found.\nInstall wgcf manually: https://github.
     com/ViRb3/wgcf/releases"
        brew install wgcf -q

    elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
        ARCH=$(uname -m)
        case $ARCH in
            x86_64)  WGCF_ARCH="amd64" ;;
            aarch64) WGCF_ARCH="arm64" ;;
            armv7l)  WGCF_ARCH="armv7" ;;
            *) die "Unsupported architecture: $ARCH.\nInstall wgcf manually: https://github.com/ViRb3/wgcf/
     releases" ;;
        esac

        echo "  Fetching latest release..."
        WGCF_VERSION=$(curl -sf https://api.github.com/repos/ViRb3/wgcf/releases/latest \
            | grep '"tag_name"' | cut -d'"' -f4)
        [[ -z "$WGCF_VERSION" ]] && die "Could not fetch wgcf version. Check your internet connection."

        sudo curl -sfL \
            "https://github.com/ViRb3/wgcf/releases/download/${WGCF_VERSION}/wgcf_${WGCF_VERSION#v}_linux_${
     WGCF_ARCH}" \
            -o /usr/local/bin/wgcf
        sudo chmod +x /usr/local/bin/wgcf
    else
```

```
        die "Unsupported OS. Install wgcf manually: https://github.com/ViRb3/wgcf/releases"
    fi
fi

ok "wgcf is ready."

#  4. WARP profile
step "Generating Cloudflare WARP profile"

if [[ -f "wgcf-profile.conf" ]]; then
    warn "wgcf-profile.conf already exists  skipping key generation."
else
    wgcf register --accept-tos 2>> output.log || die "wgcf register failed. Check your internet connection."
    wgcf generate 2>> output.log           || die "wgcf generate failed."
    ok "WARP profile generated."
fi

# Parse keys (IPv4 address only  gluetun doesn't need the IPv6 one)
PRIVATE_KEY=$(awk '/PrivateKey/{print $3}' wgcf-profile.conf)
PUBLIC_KEY=$(awk '/PublicKey/{print $3}'   wgcf-profile.conf)
ADDRESSES=$(awk '/Address/{print $3}'       wgcf-profile.conf | grep -v ':' | head -1)

[[ -z "$PRIVATE_KEY" ]] && die "Could not parse PrivateKey from wgcf-profile.conf"
[[ -z "$PUBLIC_KEY" ]]  && die "Could not parse PublicKey from wgcf-profile.conf"
[[ -z "$ADDRESSES" ]]   && die "Could not parse IPv4 Address from wgcf-profile.conf"

ok "WARP keys extracted."

#  5. Tailscale auth key
step "Tailscale authentication"

TS_AUTHKEY=""
if [[ -f ".env" ]]; then
    EXISTING_TS_KEY=$(read_env_var TS_AUTHKEY)
    if [[ -n "$EXISTING_TS_KEY" ]]; then
        TS_AUTHKEY="$EXISTING_TS_KEY"
        warn "Found existing TS_AUTHKEY in .env. Skipping prompt."
    fi
fi

if [[ -z "$TS_AUTHKEY" ]]; then
    echo ""
    echo "  Generate a key at: https://login.tailscale.com/admin/settings/keys"
    echo "   'Generate auth key'  enable Reusable if you plan to re-run setup"
    echo ""
    read -rp "  Paste your Tailscale auth key: " TS_AUTHKEY
    echo ""
fi

[[ -z "$TS_AUTHKEY" ]] && die "Tailscale auth key is required."
[[ "$TS_AUTHKEY" != tskey-auth-* ]] && warn "Key doesn't look like a Tailscale auth key  proceeding anyway."
ok "Auth key accepted."

#  6. AdGuard Home admin password
step "AdGuard Home admin account"

# Hardcoded default credentials to prevent lockout
# Username: admin
# Password: nullexit
if [[ -f ".env" ]]; then
    EXISTING_AG_PASS=$(read_env_var ADGUARD_PASSWORD)
    if [[ -n "$EXISTING_AG_PASS" ]]; then
        ADGUARD_PASSWORD="$EXISTING_AG_PASS"
        ADGUARD_PWD_ESC="$EXISTING_AG_PASS"
        ok "Found existing ADGUARD_PASSWORD in .env."
    else
        ADGUARD_PASSWORD="nullexit"
        ADGUARD_PWD_ESC="nullexit"
        ok "Default password set to 'nullexit'."
    fi
else
    ADGUARD_PASSWORD="nullexit"
    ADGUARD_PWD_ESC="nullexit"
    ok "Default password set to 'nullexit'."
fi


#  7. Write .env
step "Writing .env"

if [[ -f ".env" ]]; then
```

```
    read -rp "  .env already exists. Overwrite? [y/N]: " OVERWRITE
    if [[ "$OVERWRITE" != "y" && "$OVERWRITE" != "Y" ]]; then
        warn "Keeping existing .env."
    else
        WRITE_ENV=1
    fi
else
    WRITE_ENV=1
fi

if [[ "${WRITE_ENV:-0}" == "1" ]]; then
    # Preserve existing configuration flags if .env already exists
    EXISTING_PROFILE="medium"
    EXISTING_MSS="1120"
    EXISTING_HIJACK="true"
    EXISTING_EXIT="true"
    EXISTING_THRESH="6"
    EXISTING_KILL="true"
    EXISTING_BLOCKED="kp il"
    EXISTING_HOST_MTU="1200"
    EXISTING_SEED=""

    if [[ -f ".env" ]]; then
        [[ -n "$(read_env_var GATEWAY_RULE_PROFILE)" ]] && EXISTING_PROFILE=$(read_env_var GATEWAY_RULE_PROFILE
    )
        [[ -n "$(read_env_var GATEWAY_MSS)" ]] && EXISTING_MSS=$(read_env_var GATEWAY_MSS)
        [[ -n "$(read_env_var GATEWAY_HIJACK_HOST)" ]] && EXISTING_HIJACK=$(read_env_var GATEWAY_HIJACK_HOST)
        [[ -n "$(read_env_var GATEWAY_USE_EXIT_NODE)" ]] && EXISTING_EXIT=$(read_env_var GATEWAY_USE_EXIT_NODE)
        [[ -n "$(read_env_var WARP_FAIL_THRESHOLD)" ]] && EXISTING_THRESH=$(read_env_var WARP_FAIL_THRESHOLD)
        [[ -n "$(read_env_var KILL_SWITCH)" ]] && EXISTING_KILL=$(read_env_var KILL_SWITCH)
        [[ -n "$(read_env_var BLOCKED_COUNTRIES)" ]] && EXISTING_BLOCKED=$(read_env_var BLOCKED_COUNTRIES)
        [[ -n "$(read_env_var HOST_MTU)" ]] && EXISTING_HOST_MTU=$(read_env_var HOST_MTU)
        [[ -n "$(read_env_var NULLEXIT_SEED)" ]] && EXISTING_SEED=$(read_env_var NULLEXIT_SEED)
    fi

    # README/devref both document NULLEXIT_SEED as "generated by setup.sh"  actually
    # generate it (rather than leaving the .env.example placeholder for the user to
    # fill by hand), since crypto.sh --verify hard-fails toggle.sh's very first run
    # without one, and there's nothing else in the setup flow that would create it.
    if [[ -z "$EXISTING_SEED" ]]; then
        EXISTING_SEED=$(openssl rand -hex 32)
        ok "Generated NULLEXIT_SEED for script-integrity signing."
    fi

    cat > .env <<EOF
WIREGUARD_PRIVATE_KEY=${PRIVATE_KEY}
WIREGUARD_PUBLIC_KEY=${PUBLIC_KEY}
WIREGUARD_ADDRESSES=${ADDRESSES}
TS_AUTHKEY=${TS_AUTHKEY}
NULLEXIT_SEED=${EXISTING_SEED}
GATEWAY_RULE_PROFILE=${EXISTING_PROFILE}
# Safe MSS for double-tunneled traffic (Tailscale + WARP). Don't raise this
# devref.md 15.3.2 found 1180 (and even 1160) still too large and causes the
# exact mysterious mid-transfer stalls this value exists to prevent.
GATEWAY_MSS=${EXISTING_MSS}
# Lower the host Tailscale MTU to 1200 to prevent fragmentation when passing through the 1280-byte WARP tunnel.
HOST_MTU=${EXISTING_HOST_MTU}
# Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own
     internet is NOT hijacked).
GATEWAY_HIJACK_HOST=${EXISTING_HIJACK}
GATEWAY_USE_EXIT_NODE=${EXISTING_EXIT}
WARP_FAIL_THRESHOLD=${EXISTING_THRESH}
# PF fail-closed kill-switch. Leave true  this is the core leak guarantee.
KILL_SWITCH=${EXISTING_KILL}
ADGUARD_PASSWORD=${ADGUARD_PASSWORD}
BLOCKED_COUNTRIES="${EXISTING_BLOCKED}"
EOF
    ok ".env written."

    # Sign the core scripts against the (possibly just-generated) seed now, so the
    # very first toggle.sh run doesn't hard-fail crypto.sh --verify with no seed/no
    # .signatures and no indication that a manual `crypto.sh --sign` was needed.
    if bash scripts/crypto.sh --sign >> output.log 2>&1; then
        ok "Core scripts signed for integrity verification."
    else
        warn "Could not sign core scripts  run 'bash scripts/crypto.sh --sign' manually before toggle.sh."
    fi
fi

#  8. Prepare directories
step "Preparing AdGuard directories"
```

```
361
362 mkdir -p adguard/conf adguard/work/userfilters
363
364 # Remove empty AdGuardHome.yaml that would cause a crash loop (same logic as Toggle-Gateway scripts)
365 if [[ -f "adguard/conf/AdGuardHome.yaml" && ! -s "adguard/conf/AdGuardHome.yaml" ]]; then
366     rm "adguard/conf/AdGuardHome.yaml"
367     warn "Removed empty AdGuardHome.yaml to prevent crash loop."
368 fi
369
370 ok "Directories ready."
371
372 #  9. Compile DNS filter rules
373 step "Compiling DNS filter rules (this may take a minute)"
374
375
376 #  10. Start containers
377 step "Starting containers"
378 # Note: tailscale depends on adguardhome being healthy (port 5335 up),
379 # so Docker Compose will hold it back automatically until the wizard is done.
380 docker compose up -d
381 ok "Containers starting."
382
383 #  11. Wait for AdGuard Home setup wizard
384 step "Waiting for AdGuard Home setup wizard"
385 echo "  (up to 60 seconds...)"
386
387 MAX=60; ELAPSED=0
388 until docker exec routing-fix curl -sf http://127.0.0.1:3000 >> output.log 2>&1; do
389     sleep 2; ELAPSED=$((ELAPSED + 2))
390     [[ $ELAPSED -ge $MAX ]] && die "AdGuard Home didn't come up.\nDebug with: docker compose logs adguardhome"
391 done
392
393 ok "AdGuard Home is up."
394
395 #  12. Configure AdGuard Home via API
396 step "Configuring AdGuard Home"
397
398 # Run all API calls in a single docker exec sh -c so the session cookie file
399 # is shared across calls without leaving the container.
400 docker exec routing-fix sh -c "
401   set -e
402
403   # Step 1: Complete the setup wizard (sets admin creds + DNS port to 5335)
404   curl -sf -X POST http://127.0.0.1:3000/control/install/configure \
405     -H 'Content-Type: application/json' \
406     -d '{\"web\":{\"ip\":\"0.0.0.0\",\"port\":3000},\"dns\":{\"ip\":\"0.0.0.0\",\"port\":5335},\"username\":\"
      admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
407     >/dev/null || true
408
409   # Step 2: Wait for AdGuard to restart after wizard
410   sleep 4
411   WAITED=0
412   until curl -sf http://127.0.0.1:3000 >/dev/null 2>&1; do
413     sleep 2; WAITED=\$((WAITED + 2))
414     [ \$WAITED -ge 30 ] && { echo 'AdGuard did not restart in time'; exit 1; }
415   done
416
417   # Step 3: Login to get session cookie
418   curl -sf -X POST http://127.0.0.1:3000/control/login \
419     -H 'Content-Type: application/json' \
420     -d '{\"name\":\"admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
421     -c /tmp/agh.cookie >/dev/null
422
423   # Step 4: Point upstream DNS back through the WARP tunnel
424   curl -sf -X POST http://127.0.0.1:3000/control/dns_config \
425     -H 'Content-Type: application/json' \
426     -b /tmp/agh.cookie \
427     -d '{\"upstream_dns\":[\"127.0.0.1:53\",\"[/onion/]127.0.0.1:5353\"],\"bootstrap_dns\":[\"1.1.1.1\",\"
      8.8.8.8\"],\"upstream_mode\":\"load_balance\"}' \
428     >/dev/null
429
430   # Step 5: Register the compiled rules file as a filter list
431   curl -sf -X POST http://127.0.0.1:3000/control/filtering/add_url \
432     -H 'Content-Type: application/json' \
433     -b /tmp/agh.cookie \
434     -d '{\"name\":\"Custom Compiled Rules\",\"url\":\"/opt/adguardhome/work/userfilters/compiled_rules.txt\"}'
      \
435     >/dev/null || true
436
437   # Step 6: Force a filter refresh to load the rules immediately
438   curl -sf -X POST http://127.0.0.1:3000/control/filtering/refresh \
```

```
    -H 'Content-Type: application/json' \
    -b /tmp/agh.cookie \
    -d '{\"whitelist\":false}' \
    >/dev/null || true

  rm -f /tmp/agh.cookie
"

ok "AdGuard Home configured."

#  13. Wait for Tailscale
step "Waiting for Tailscale to authenticate"
echo "  (Tailscale only starts once AdGuard's healthcheck passes  up to 90s)"

MAX=90; ELAPSED=0; TS_IP=""
until [[ -n "$TS_IP" ]]; do
    TS_IP=$(docker exec tailscale tailscale ip -4 2>> output.log || true)
    if [[ -z "$TS_IP" ]]; then
        sleep 3; ELAPSED=$((ELAPSED + 3))
        if [[ $ELAPSED -ge $MAX ]]; then
            warn "Tailscale hasn't authenticated yet."
            warn "Once it does, re-run the toggle script  it will resolve the IP dynamically."
            break
        fi
    fi
done

if [[ -n "$TS_IP" ]]; then
    ok "Tailscale IP: ${TS_IP} (resolved dynamically by the toggle script on each startup)"
fi
```

53 scripts/sweep.sh

```
#!/bin/bash
# scripts/sweep.sh  nullexit gateway health sweep
#
# Automates the 7-check post-restart health sweep (sweep.md 1 + tailnet ping
# + throughput). It answers one question: "after a toggle/restart/wake, is
# the gateway actually healthy end-to-end  no leak, direct P2P, DNS
# filtering live, kill-switch armed, and fast enough to actually use?" It is
# READ-ONLY with respect to gateway state: it never mutates routing, PF, DNS,
# or containers (the throughput check only pulls a test file over curl  the
# same kind of traffic any app already sends, nothing is exec'd into or
# installed in the container), so it is safe to run against a live gateway at
# any time (including from a remote session  see the caveat in agent.md L16).
#
# The 7 checks (1-5 mirror sweep.md 1):
#   1. Containers           warp, adguardhome, tailscale, tor, routing-fix up/healthy
#   2. Double-tunnel/leak   container WARP exit IP == host exit IP, both warp=on
#   3. Hostcontainer path   Tailscale exit-node is a DIRECT P2P conn, not DERP relay
#   4. AdGuard filtering     compiled rule counts present + live DNS interception
#   5. PF kill-switch        pfctl Enabled + anchor com.apple/nullexit carries rules
#   6. Tailnet reachability  every ONLINE Tailscale peer answers a tailscale ping
#      (DERP-relayed peers are flagged as throughput-degraded, not a failure)
#   7. Throughput            host path vs. WARP-only baseline (via the container's
#      SOCKS5 proxy, no exec/install inside the container). Catches "slow AND
#      dropping mid-transfer" that a one-shot ping/latency check can't see.
#      Real per-device (phone) throughput can't be measured without an agent
#      running on the device, so mesh peers get a path/latency flag instead
#      (see check 6).
#
# A timestamped report is written to the project root (one file per run, so
# runs can be diffed). stdout carries a coloured summary + a final PASS/FAIL
# verdict; exit code is 0 when every check passes, 1 otherwise (so it can gate
# CI / launchd / a post-restart hook).
#
# Usage:
#   bash scripts/sweep.sh            # run the sweep, write report, print verdict
#   bash scripts/sweep.sh --quiet    # only the final verdict line + non-zero on FAIL
#   bash scripts/sweep.sh --help     # this message

set -uo pipefail
# NOTE: errexit is intentionally OFF. Several probes are EXPECTED to fail on an
# unhealthy gateway and that failure IS the diagnostic signal  we record each
# exit status rather than letting one abort the whole sweep.

```

```
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
source "$SCRIPT_DIR/common.sh"
LOG_FILE="$PROJECT_ROOT/output.log"
REPORT_FILE="$PROJECT_ROOT/sweep-$TIMESTAMP.txt"

#  Argument parsing
QUIET=false
case "${1:-}" in
  --quiet|-q) QUIET=true ;;
  --help|-h)  sed -n '2,27p' "$0"; exit 0 ;;
  "")         : ;;
  *)          echo "Unknown argument: $1" >&2
              echo "Run with --help for usage." >&2
              exit 2 ;;
esac

#  Colours (auto-disabled when stdout is not a tty, e.g. piped)
if [ -t 1 ]; then
  RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
  BOLD='\033[1m'; NC='\033[0m'
else
  RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
fi

#  Homebrew on PATH (same as toggle.sh / recover.sh / diagnose-host-leak)
export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"

#  Output helpers  write to BOTH stdout and the report file
exec 3>> "$REPORT_FILE" || {
  printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
  exit 1
}
emit()    { [ "$QUIET" = true ] || printf '%b\n' "$1"; printf '%b\n' "$1" | sed 's/\x1b\[[0-9;]*m//g' >&3; }
section() { emit ""; emit "${BOLD} $1 ${NC}"; }
line()    { emit "  $1"; }
ok()      { line "${GREEN}${NC} $1"; }
warn()    { line "${YELLOW}${NC} $1"; }
fail()    { line "${RED}${NC} $1"; }

# common.sh redefines run_with_timeout; it is a pure-bash macOS-safe timeout
# (GNU `timeout` is absent on stock macOS). We reuse the sourced version.

#  Result accounting
# Each check appends "PASS"/"WARN"/"FAIL" so the verdict is a pure roll-up.
declare -a RESULTS=()
record() { RESULTS+=("$1|$2"); }   # status|label

#  Header
if [ "$QUIET" != true ]; then
  emit ""
  emit ""
  emit " n u l l e x i t   h e a l t h   s w e e p"
  emit ""
  emit " Timestamp (UTC): $TIMESTAMP"
  emit " Project root:    $PROJECT_ROOT"
  emit " Report file:     $REPORT_FILE"
  emit ""
fi

GATEWAY_TS_IP="$(read_adguard_ip)"   # gateway Tailscale IP from .gateway_ip

#
section "1/7  Containers"
#
# All five services must be running; warp + adguardhome additionally expose a
# healthcheck and must be 'healthy'. `docker compose ps` is the source of truth.
EXPECTED_SERVICES="warp adguardhome tailscale tor routing-fix"
if ! command -v docker >/dev/null 2>&1; then
  fail "docker CLI not on PATH  cannot inspect containers"
  record FAIL "containers"
else
  PS_OUT=$(run_with_timeout 20 docker compose ps 2>>"$LOG_FILE")
  missing=""; unhealthy=""
  for svc in $EXPECTED_SERVICES; do
    row=$(printf '%s\n' "$PS_OUT" | awk -v s="$svc" '$1==s || $0 ~ ("[[:space:]]"s"[[:space:]]")')
    if [ -z "$row" ] || ! printf '%s' "$row" | grep -qiE 'up|running'; then
      missing="$missing $svc"
    elif printf '%s' "$row" | grep -qi 'unhealthy'; then
```

```
      unhealthy="$unhealthy $svc"
    fi
  done
  if [ -z "$missing" ] && [ -z "$unhealthy" ]; then
    ok "all 5 services up (warp/adguardhome healthy, tailscale/tor/routing-fix running)"
    record PASS "containers"
  else
    [ -n "$missing" ]   && fail "not up/running:$missing"
    [ -n "$unhealthy" ] && fail "reporting unhealthy:$unhealthy"
    record FAIL "containers"
  fi
  line " docker compose ps "
  printf '%s\n' "$PS_OUT" | awk '{print "    "$0}' >&3
  [ "$QUIET" = true ] || printf '%s\n' "$PS_OUT" | awk '{print "    "$0}'
fi

#
section "2/7  Double-tunnel / IP-leak"
#
# The container's WARP egress and the host's egress must be the SAME public IP
# with warp=on for both. Equal IPs  all host traffic is exiting through the
# same WARP tunnel the container uses (no split / raw-ISP leak).
# The warp container (gluetun) ships wget, not curl  fetch the trace from
# inside its netns with whichever client is present.
CTR_TRACE=$(run_with_timeout 12 docker compose exec -T warp sh -c \
            'curl -s --max-time 8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
             || wget -qO- --timeout=8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null' \
            2>>"$LOG_FILE")
HOST_TRACE=$(run_with_timeout 12 curl -s --max-time 8 \
            https://www.cloudflare.com/cdn-cgi/trace 2>>"$LOG_FILE")
CTR_IP=$(printf '%s' "$CTR_TRACE"  | awk -F= '/^ip=/{print $2;exit}')
CTR_WARP=$(printf '%s' "$CTR_TRACE" | awk -F= '/^warp=/{print $2;exit}')
HOST_IP=$(printf '%s' "$HOST_TRACE" | awk -F= '/^ip=/{print $2;exit}')
HOST_WARP=$(printf '%s' "$HOST_TRACE"| awk -F= '/^warp=/{print $2;exit}')
line "container exit: ip=${CTR_IP:-?} warp=${CTR_WARP:-?}"
line "host exit:      ip=${HOST_IP:-?} warp=${HOST_WARP:-?}"
if [ -z "$CTR_IP" ] || [ -z "$HOST_IP" ]; then
  fail "could not obtain both egress IPs (container or host trace failed)"
  record FAIL "leak"
elif [ "$CTR_IP" = "$HOST_IP" ] && [ "$HOST_WARP" = "on" ]; then
  ok "no leak  host egress == container egress ($HOST_IP), warp=on"
  record PASS "leak"
elif [ "$HOST_WARP" != "on" ]; then
  fail "host warp=$HOST_WARP  host traffic is NOT going through WARP (LEAK)"
  record FAIL "leak"
else
  fail "egress mismatch: host $HOST_IP != container $CTR_IP  split/leak"
  record FAIL "leak"
fi

#
section "3/7  Host  container path (direct P2P, not DERP relay)"
#
# `tailscale status` annotates each peer with 'direct <ip:port>' or
# 'relay "<derp>"'. The gateway peer should be DIRECT  a DERP relay hop adds
# latency and signals NAT traversal failure between host and gateway.
if ! command -v tailscale >/dev/null 2>&1; then
  fail "tailscale CLI not on PATH"
  record FAIL "p2p"
else
  TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
  PEER_LINE=$(printf '%s\n' "$TS_OUT" | grep -E "(^|[[:space:]])$GATEWAY_TS_IP([[:space:]]|$)" | head -1)
  if [ -z "$PEER_LINE" ]; then
    warn "gateway $GATEWAY_TS_IP not found in host peer list"
    line "$(printf '%s\n' "$TS_OUT" | head -5 | awk '{print "    "$0}')"
    record WARN "p2p"
  elif printf '%s' "$PEER_LINE" | grep -q 'direct '; then
    conn=$(printf '%s' "$PEER_LINE" | grep -oE 'direct [0-9.]+:[0-9]+')
    ok "direct P2P to gateway  $conn (no DERP relay)"
    record PASS "p2p"
  elif printf '%s' "$PEER_LINE" | grep -q 'relay '; then
    relay=$(printf '%s' "$PEER_LINE" | grep -oE 'relay "[^"]+"')
    warn "gateway reached via DERP $relay (not direct)  NAT traversal degraded"
    record WARN "p2p"
  else
    warn "gateway peer state indeterminate (idle/no active conn)"
    line "    $PEER_LINE"
    record WARN "p2p"
  fi
fi

```

```
#
section "4/7  AdGuard filtering"
#
# Two signals: (a) the compiled ruleset exists with a non-trivial block count,
# (b) the gateway resolver actually answers a query (live interception).
COMPILED="adguard/work/userfilters/compiled_rules.txt"
if [ -f "$COMPILED" ]; then
  TOTAL=$(grep -m1 '! Total Block Rules:'  "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
  NATIVE=$(grep -m1 '! Native AdGuard Rules:' "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
  if [ -n "$TOTAL" ] && [ "$TOTAL" -gt 0 ] 2>/dev/null; then
    ok "compiled ruleset present  Total Block Rules: $TOTAL / Native AdGuard: ${NATIVE:-?}"
    RULES_OK=true
  else
    warn "compiled_rules.txt present but block-rule count unreadable"
    RULES_OK=false
  fi
else
  warn "compiled_rules.txt missing (rule-compile may not have run)"
  RULES_OK=false
fi
# Live interception: the gateway resolver must answer (dig via .gateway_ip).
DNS_OK=false
if command -v dig >/dev/null 2>&1 && [ -n "$GATEWAY_TS_IP" ]; then
  DIG_OUT=$(run_with_timeout 8 dig +time=3 +tries=1 @"$GATEWAY_TS_IP" google.com A 2>>"$LOG_FILE")
  if printf '%s' "$DIG_OUT" | grep -qE 'ANSWER SECTION|status: NOERROR'; then
    ok "live DNS interception  gateway $GATEWAY_TS_IP resolved google.com"
    DNS_OK=true
  else
    warn "gateway $GATEWAY_TS_IP did not answer a test query"
  fi
else
  warn "skipped live-resolution probe (dig missing or gateway IP unknown)"
fi
if [ "$RULES_OK" = true ] && [ "$DNS_OK" = true ]; then
  record PASS "adguard"
elif [ "$RULES_OK" = true ] || [ "$DNS_OK" = true ]; then
  record WARN "adguard"
else
  record FAIL "adguard"
fi

#
section "5/7  PF kill-switch"
#
# PF must be Enabled AND the com.apple/nullexit anchor must carry rules.
# An empty/absent anchor means the fail-closed lane is not actually armed.
PF_STATUS=$(run_with_timeout 6 sudo -n pfctl -s info 2>>"$LOG_FILE" | awk '/^Status:/{print $2;exit}')
ANCHOR_RULES=$(run_with_timeout 6 sudo -n pfctl -a com.apple/nullexit -sr 2>>"$LOG_FILE")
if [ "$PF_STATUS" = "Enabled" ] && [ -n "$ANCHOR_RULES" ]; then
  RULE_N=$(printf '%s\n' "$ANCHOR_RULES" | grep -cvE '^\s*$')
  ok "PF Enabled + anchor com.apple/nullexit armed ($RULE_N rules)"
  line "$(printf '%s\n' "$ANCHOR_RULES" | head -4 | awk '{print "    "$0}')"
  record PASS "pf"
elif [ "$PF_STATUS" = "Enabled" ]; then
  fail "PF Enabled but anchor com.apple/nullexit has NO rules  kill-switch not armed"
  record FAIL "pf"
elif [ -z "$PF_STATUS" ]; then
  warn "could not read PF status (needs passwordless sudo for 'pfctl -s info')"
  record WARN "pf"
else
  fail "PF Status: $PF_STATUS  kill-switch disabled"
  record FAIL "pf"
fi

#
section "6/7  Tailnet reachability (burst ping  loss % + jitter, all online devices)"
#
# Every peer the coordination server reports as ONLINE should answer a
# `tailscale ping` (TSMP probes the WireGuard data path; ICMP is avoided since
# phones/laptops often drop it). Offline peers and this host are skipped.
# An online-but-unreachable peer means control plane and data path disagree.
#
# A single ping only proves "reachable right now"  it can't see loss or
# jitter, which is what actually determines whether a call/video/QUIC session
# on that device holds up. There's no way to get a real bits-per-second number
# for a phone without something running on it to receive/echo data (neither
# phone has that), but TSMP ping is answered by the Tailscale client itself
# with zero cooperation needed from the device  the one channel that works
# identically over DERP or direct P2P. So: send a BURST instead of one ping,
# and report loss % + RTT spread per peer, the closest honest signal to
# "connection quality" available without an app/listener on the phone.
```

```
PING_BURST="${SWEEP_PING_BURST:-8}"
if ! command -v tailscale >/dev/null 2>&1; then
  fail "tailscale CLI not on PATH"
  record FAIL "tailnet"
else
  SELF_TS_IP=$(run_with_timeout 6 tailscale ip -4 2>>"$LOG_FILE")
  [ -n "${TS_OUT:-}" ] || TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
  # Online peers = rows starting with a Tailscale IP, minus self and 'offline'.
  # (A '-' status column means online-but-idle, so it is kept.)
  PEERS=$(printf '%s\n' "$TS_OUT" | awk -v self="$SELF_TS_IP" \
          '$1 ~ /^100\./ && $1 != self && $0 !~ /(^|[[:space:]])offline(,|;|[[:space:]]|$)/ {print $1, $2}')
  if [ -z "$PEERS" ]; then
    warn "no online peers to ping (everything else in the tailnet is offline)"
    record WARN "tailnet"
  else
    unreached=0; lossy=0
    while read -r peer_ip peer_name; do
      PONGS=$(run_with_timeout $((PING_BURST + 8)) tailscale ping --timeout=3s --c "$PING_BURST" --until-direct=false "$peer_ip" 2>>"$LOG_FILE")
      n_ok=$(printf '%s\n' "$PONGS" | grep -c '^pong')
      if [ "$n_ok" -eq 0 ]; then
        fail "$peer_name ($peer_ip)  no pong in $PING_BURST attempts (online per control plane, data path dead)"
        unreached=$((unreached+1))
        continue
      fi
      path=$(printf '%s\n' "$PONGS" | grep -m1 '^pong' | sed -E 's/^pong from [^ ]+ \([^)]+\) via ([^ ]+).*/\1/')
      loss_pct=$(( (PING_BURST - n_ok) * 100 / PING_BURST ))
      read -r avg_ms min_ms max_ms <<< "$(printf '%s\n' "$PONGS" | grep '^pong' | grep -oE '[0-9]+ms' | tr -d 'ms' | awk '
        { if (NR==1 || $1<min) min=$1; if (NR==1 || $1>max) max=$1; sum+=$1; n++ }
        END { if (n>0) printf "%.0f %d %d", sum/n, min, max }')"
      jitter_ms=$((max_ms - min_ms))
      relay_note=""
      case "$path" in
        DERP*) relay_note="  relayed, expect degraded throughput (see check 7)" ;;
      esac
      if [ "$loss_pct" -eq 0 ]; then
        ok "$peer_name ($peer_ip)  $n_ok/$PING_BURST pong (0% loss), RTT ${min_ms}-${max_ms}ms avg ${avg_ms}ms (jitter ${jitter_ms}ms) via ${path}${relay_note}"
      elif [ "$loss_pct" -lt 40 ]; then
        warn "$peer_name ($peer_ip)  $n_ok/$PING_BURST pong (${loss_pct}% loss), RTT ${min_ms}-${max_ms}ms avg ${avg_ms}ms via ${path}${relay_note}"
        lossy=$((lossy+1))
      else
        fail "$peer_name ($peer_ip)  $n_ok/$PING_BURST pong (${loss_pct}% loss, unstable) via ${path}${relay_note}"
        unreached=$((unreached+1))
      fi
    done <<< "$PEERS"
    n_peers=$(printf '%s\n' "$PEERS" | wc -l | tr -d '[:space:]')
    if [ "$unreached" -eq 0 ] && [ "$lossy" -eq 0 ]; then
      record PASS "tailnet"
    elif [ "$unreached" -lt "$n_peers" ]; then
      record WARN "tailnet"
    else
      record FAIL "tailnet"
    fi
  fi
fi

#
section "7/7  Throughput (host path vs. WARP-only baseline)"
#
# A single ping only proves the path is UP, not that it stays up under real
# load  devref.md 15.3.1 documents sustained-transfer bufferbloat/reset
# behavior a quick latency check can't see. Pulls a real test file twice:
# once over the host's normal route (full double-tunnel: Tailscale exit-node
# wrap + WARP) and once straight through the warp container's own SOCKS5
# proxy (WARP only, no Tailscale wrap)  port read from the live
# /tmp/nullexit-ports.env, so this never execs into or installs anything in
# the container (sweep stays read-only w.r.t. gateway state). Comparing the
# two shows whether slowness/resets come from WARP itself or from the extra
# Tailscale-wrap hop. Skip with SWEEP_SKIP_THROUGHPUT=true in .env.
#
# Target: a fast, high-bandwidth CDN, not a slow/distant one. speedtest.net-style
# test hosts (e.g. speedtest.tele2.net) were tried first and turned out to be the
# bottleneck themselves  even RAW, gateway-off ISP speed to tele2.net was ~3.2
# Mbps, vs ~32 Mbps raw to a fast CDN  so a slow test target would have silently
# blamed nullexit for a benchmark artifact. Cloudflare's own speed-test endpoint
```

```
# 403s WARP egress IPs specifically (abuse-prevention on shared/CGNAT IPs), so
# Hetzner (also fast, does not flag WARP IPs) is used instead.
THROUGHPUT_SKIP=$(read_env_var "SWEEP_SKIP_THROUGHPUT" | tr '[:upper:]' '[:lower:]')
if [ "$THROUGHPUT_SKIP" = "true" ]; then
  warn "skipped (SWEEP_SKIP_THROUGHPUT=true)"
  record WARN "throughput"
else
  TP_URL="${SWEEP_THROUGHPUT_URL:-https://ash-speed.hetzner.com/100MB.bin}"
  TP_TIMEOUT="${SWEEP_THROUGHPUT_TIMEOUT:-30}"

  measure() {   # $1 = extra curl args (word-split; may be empty)
    # curl's own --max-time must win the race against run_with_timeout's outer
    # SIGTERM watcher  a process killed by signal never reaches its own -w
    # output, so if both timers were equal, the watcher could kill curl right
    # before it prints stats, and every leg would misreport as "no response".
    # The outer timeout is a dead-man's-switch (+5s), not the real budget.
    run_with_timeout $((TP_TIMEOUT + 5)) curl -sS --max-time "$TP_TIMEOUT" $1 -o /dev/null \
      -w 'code=%{http_code} size=%{size_download} time=%{time_total} bps=%{speed_download}' \
      "$TP_URL" 2>>"$LOG_FILE"
  }
  leg_verdict() {   # $1=curl_rc $2=http_code -> ok|stall|blocked|unknown
    [ -z "$2" ] && { echo unknown; return; }
    [ "$2" != "200" ] && { echo blocked; return; }
    case "$1" in
      0|28) echo ok ;;     # 0=completed; 28=hit our own --max-time while still healthy
      56)   echo stall ;;  # curl's actual "Recv failure: Connection reset by peer"
      *)    echo unknown ;;
    esac
  }
  to_mbps() { awk -v b="${1:-}" 'BEGIN{ if (b=="") print "?"; else printf "%.1f", (b*8)/1000000 }'; }

  HOST_TP=$(measure ""); HOST_RC=$?
  HOST_CODE=$(printf '%s' "$HOST_TP" | grep -oE 'code=[0-9]+' | cut -d= -f2)
  HOST_SIZE=$(printf '%s' "$HOST_TP" | grep -oE 'size=[0-9]+' | cut -d= -f2)
  HOST_TIME=$(printf '%s' "$HOST_TP" | grep -oE 'time=[0-9.]+' | cut -d= -f2)
  HOST_BPS=$(printf '%s' "$HOST_TP"  | grep -oE 'bps=[0-9.]+'  | cut -d= -f2)
  HOST_VERDICT=$(leg_verdict "$HOST_RC" "$HOST_CODE")
  HOST_MBPS=$(to_mbps "$HOST_BPS")
  case "$HOST_VERDICT" in
    ok)      ok   "host path: ${HOST_MBPS} Mbps (${HOST_SIZE:-0} bytes in ${HOST_TIME:-?}s)" ;;
    stall)   warn "host path: connection reset mid-transfer after ${HOST_SIZE:-0} bytes  sustained-load drop,
     see devref.md 15.3.1" ;;
    blocked) warn "host path: test endpoint returned HTTP ${HOST_CODE:-?} (egress IP likely flagged)
     inconclusive, not a gateway fault" ;;
    unknown) warn "host path: no response (timeout after ${TP_TIMEOUT}s)" ;;
  esac

  SOCKS_PROXY_PORT=""
  [ -f /tmp/nullexit-ports.env ] && SOCKS_PROXY_PORT=$(grep -oE 'SOCKS_PROXY_PORT=[0-9]+' /tmp/nullexit-ports.
     env | cut -d= -f2)
  if [ -z "$SOCKS_PROXY_PORT" ]; then
    warn "WARP-only baseline: SOCKS_PROXY_PORT unknown (no /tmp/nullexit-ports.env  gateway not started via
     toggle.sh?)"
    CTR_VERDICT=unknown
  else
    CTR_TP=$(measure "--socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT"); CTR_RC=$?
    CTR_CODE=$(printf '%s' "$CTR_TP" | grep -oE 'code=[0-9]+' | cut -d= -f2)
    CTR_SIZE=$(printf '%s' "$CTR_TP" | grep -oE 'size=[0-9]+' | cut -d= -f2)
    CTR_TIME=$(printf '%s' "$CTR_TP" | grep -oE 'time=[0-9.]+' | cut -d= -f2)
    CTR_BPS=$(printf '%s' "$CTR_TP"  | grep -oE 'bps=[0-9.]+'  | cut -d= -f2)
    CTR_VERDICT=$(leg_verdict "$CTR_RC" "$CTR_CODE")
    CTR_MBPS=$(to_mbps "$CTR_BPS")
    case "$CTR_VERDICT" in
      ok)      ok   "WARP-only baseline: ${CTR_MBPS} Mbps (${CTR_SIZE:-0} bytes in ${CTR_TIME:-?}s)" ;;
      stall)   warn "WARP-only baseline: connection reset mid-transfer after ${CTR_SIZE:-0} bytes  bug is in
     WARP itself, not the Tailscale wrap" ;;
      blocked) warn "WARP-only baseline: test endpoint returned HTTP ${CTR_CODE:-?}  inconclusive" ;;
      unknown) warn "WARP-only baseline: no response (timeout after ${TP_TIMEOUT}s)" ;;
    esac
  fi

  if [ "$HOST_VERDICT" = "ok" ] && [ "$CTR_VERDICT" = "ok" ]; then
    if [ "$HOST_MBPS" != "?" ] && [ "$CTR_MBPS" != "?" ]; then
      DELTA=$(awk -v h="$HOST_MBPS" -v c="$CTR_MBPS" 'BEGIN{ if (c<=0) print 0; else printf "%.0f", 100*(c-h)/c
      }')
      [ "$DELTA" -ge 30 ] 2>/dev/null && warn "host is ${DELTA}% slower than the WARP-only baseline  the
     Tailscale exit-node wrap is the added cost here, not WARP"
    fi
    record PASS "throughput"
  elif [ "$HOST_VERDICT" = "stall" ] || [ "$CTR_VERDICT" = "stall" ]; then
    record WARN "throughput"
```

```
  elif [ "$HOST_VERDICT" = "unknown" ] && [ "$CTR_VERDICT" = "unknown" ]; then
    record FAIL "throughput"
  else
    record WARN "throughput"
  fi
fi

#
# Cover traffic (opt-in)  informational only, NOT a scored check.
# Shown solely when NOISE_ENABLED=true so the sweep stays read-only and the
# 6-check verdict is unchanged. Reports whether the padding daemon is alive.
NOISE_ON=$(read_env_var "NOISE_ENABLED" | tr '[:upper:]' '[:lower:]')
if [ "$NOISE_ON" = "true" ]; then
  section "Cover traffic (opt-in, informational)"
  NOISE_PID_FILE="/tmp/nullexit-noise.pid"
  npid="$(cat "$NOISE_PID_FILE" 2>/dev/null)"
  if [ -n "$npid" ] && kill -0 "$npid" 2>/dev/null; then
    ok "dummy-padding daemon running (pid $npid)  see 'bash scripts/noise.sh status'"
  else
    warn "NOISE_ENABLED=true but no padding daemon running  start it: bash scripts/noise.sh start"
  fi
fi

#
# DNS anomaly scan (informational, read-only, NOT a scored check  the 6-check
# verdict is unchanged). It only READS the AdGuard querylog; never blocks DNS.
# It FAILS LOUD: a broken detector (crash, corrupt model, empty log, failing
# self-test) shows a red  with the reason  it is NEVER silently reported as
# clean, and an untrained detector is announced rather than skipped in silence.
QUERYLOG="$PROJECT_ROOT/adguard/work/data/querylog.json"
DNS_TOOL="$SCRIPT_DIR/dns_anomaly_detector.py"
if command -v python3 >/dev/null 2>&1 && [ -f "$DNS_TOOL" ]; then
  section "DNS anomaly scan (informational)"
  # 1) Prove the detection logic itself works before trusting any "clean".
  ST_ERR=$(python3 "$DNS_TOOL" selftest 2>&1 >/dev/null); ST_RC=$?
  if [ "$ST_RC" -ne 0 ]; then
    fail "DNS detector SELF-TEST FAILED (rc=$ST_RC)  detector is broken, do NOT trust results: ${ST_ERR:-see
     stderr}"
  elif [ ! -f "$PROJECT_ROOT/.dns_baseline.json" ]; then
    warn "DNS detector NOT trained (no .dns_baseline.json)  not protecting; run: python3 scripts/
     dns_anomaly_detector.py learn"
  elif [ ! -f "$QUERYLOG" ]; then
    warn "DNS detector trained but querylog absent ($QUERYLOG)  nothing to scan"
  else
    # 2) Real scan. Capture stderr + exit code; do NOT swallow failures.
    DNS_ERR=$(python3 "$DNS_TOOL" scan --quiet 2>&1 >/tmp/nullexit-dns-scan.out); DNS_RC=$?
    DNS_OUT=$(cat /tmp/nullexit-dns-scan.out 2>/dev/null); rm -f /tmp/nullexit-dns-scan.out
    case "$DNS_RC" in
      0) ok   "no DNS-tunneling / DGA signatures  $DNS_OUT" ;;
      1) warn "$DNS_OUT  review: python3 scripts/dns_anomaly_detector.py scan" ;;
      *) fail "DNS scan FAILED to run (rc=$DNS_RC)  NOT a clean result: ${DNS_ERR:-unknown error}" ;;
    esac
  fi
fi

#
section "VERDICT"
#
pass=0; warnc=0; failc=0
for r in "${RESULTS[@]}"; do
  case "${r%%|*}" in PASS) pass=$((pass+1));; WARN) warnc=$((warnc+1));; FAIL) failc=$((failc+1));; esac
done
for r in "${RESULTS[@]}"; do
  st="${r%%|*}"; lbl="${r##*|}"
  case "$st" in
    PASS) ok   "$lbl" ;;
    WARN) warn "$lbl" ;;
    FAIL) fail "$lbl" ;;
  esac
done
emit ""
if [ "$failc" -eq 0 ] && [ "$warnc" -eq 0 ]; then
  emit "${GREEN}${BOLD}SWEEP PASS  gateway healthy (${pass}/${#RESULTS[@]} checks green)${NC}"
  EXIT=0
elif [ "$failc" -eq 0 ]; then
  emit "${YELLOW}${BOLD}SWEEP PASS with warnings  ${pass} pass, ${warnc} warn, 0 fail${NC}"
  EXIT=0
else
  emit "${RED}${BOLD}SWEEP FAIL  ${pass} pass, ${warnc} warn, ${failc} fail${NC}"
  EXIT=1
fi
```

```
emit ""
[ "$QUIET" = true ] || emit "Full report: $REPORT_FILE"
exit "$EXIT"
```

54 scripts/taildrop-watch.sh

```
#!/bin/bash
# scripts/taildrop-watch.sh
# launchd entry point for ~/Library/LaunchAgents/com.nullexit.taildrop-watch.plist
# (label com.nullexit.taildrop-watch).
#
# Homebrew's tailscaled CLI (unlike the Tailscale.app GUI) does not auto-drain
# incoming Taildrop files  they sit in an internal inbox until something calls
# `tailscale file get`. This blocks in --wait --loop mode so every incoming
# file moves into ~/Downloads within moments of arriving instead of sitting
# invisible until someone remembers to pull it manually.

set -euo pipefail
export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"

exec tailscale file get --wait --loop --conflict=rename "$HOME/Downloads"
```

55 scripts/unlock-files.sh

```
#!/bin/bash
# Resolve project root relative to this script's location, so the script
# works regardless of the caller's CWD.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"

echo "Unlocking .env..."
if [ -f "$PROJECT_ROOT/.env" ]; then
  sudo cat "$PROJECT_ROOT/.env" > "$PROJECT_ROOT/.env.tmp" && mv "$PROJECT_ROOT/.env.tmp" "$PROJECT_ROOT/.env"
  echo " .env unlocked"
fi

echo "Unlocking AdGuard cache files..."
for f in "$PROJECT_ROOT"/adguard/work/userfilters/compiled_rules.txt \
         "$PROJECT_ROOT"/adguard/work/userfilters/ip_blocklist.ipset \
         "$PROJECT_ROOT"/adguard/work/userfilters/cache/*.txt \
         "$PROJECT_ROOT"/adguard/work/userfilters/cache/ip/*.txt \
         "$PROJECT_ROOT"/adguard/work/data/filters/*.txt; do
  if [ -f "$f" ]; then
    cat "$f" > "$f.tmp" && mv "$f.tmp" "$f"
    echo " Unlocked $f"
  fi
done

echo ""
echo "All locked files have been successfully bypassed via atomic rename!"
```

56 scripts/wifi-rejoin.sh

```
#!/usr/bin/env bash
# nullexit  Wi-Fi keepalive. Re-associates ONLY to networks you've remembered,
# never any other (no public-AP hopping). No Location permission, no root plist
# reads, no reading the current SSID.
#
# Why "remembered" and not "current": macOS hides the live SSID from every
# script unless the process holds Location Services access, so nothing can
# auto-detect which network you're on. Joining a network *by name* is not
# gated, so you remember your network(s) once and we rejoin the first
# remembered one that's in range.
#
#   set <SSID>      remember a network (safe to run repeatedly)
#   forget <SSID>   stop remembering it
#   list            show remembered networks
#   reset           (internal) clear the drop-timer once the link is back
```

```
#   rejoin          (internal, called by the WARP watcher) re-associate a
#                   remembered network after a grace period, with backoff
set -u

DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
ROOT="$(cd "$DIR/.." && pwd)"
STATE="$ROOT/.last_wifi_ssid"          # remembered SSIDs, one per line (gitignored)
LOG="$ROOT/output.log"
DOWN_MARK="/tmp/nullexit-wifi-downsince"
JOIN_MARK="/tmp/nullexit-wifi-lastjoin"
TRY_MARK="/tmp/nullexit-wifi-tries"
GAVEUP_MARK="/tmp/nullexit-wifi-gaveup"

GRACE=15        # let DHCP self-heal before the first rejoin attempt
BACKOFF=30      # base seconds between rejoin attempts (grows per attempt)
MAX_TRIES=6     # stop re-associating after this many tries per down-episode

ts()  { date -u +%FT%TZ; }
log() { echo "[$(ts)] wifi-rejoin: $*" >> "$LOG"; }

wifi_iface() {
  local i
  i=$(networksetup -listallhardwareports 2>/dev/null \
      | awk '/Hardware Port: Wi-Fi/{getline; print $2; exit}')
  [ -z "$i" ] && i=en0
  printf '%s' "$i"
}

case "${1:-}" in
  set)
    ssid="${2:-}"; [ -z "$ssid" ] && { echo "usage: ${0##*/} set <SSID>" >&2; exit 2; }
    touch "$STATE"
    if grep -qxF "$ssid" "$STATE" 2>/dev/null; then
      echo "already remembered: $ssid"
    else
      printf '%s\n' "$ssid" >> "$STATE"
      echo "remembered: $ssid"
    fi
    ;;

  forget)
    ssid="${2:-}"; [ -z "$ssid" ] && { echo "usage: ${0##*/} forget <SSID>" >&2; exit 2; }
    if [ -f "$STATE" ] && grep -qxF "$ssid" "$STATE"; then
      grep -vxF "$ssid" "$STATE" > "$STATE.tmp" && mv "$STATE.tmp" "$STATE"
      echo "forgotten: $ssid"
    else
      echo "not remembered: $ssid"
    fi
    ;;

  list)
    if [ -s "$STATE" ]; then cat "$STATE"; else echo "(none  run: ${0##*/} set <SSID>)"; fi
    ;;

  reset)
    rm -f "$DOWN_MARK" "$JOIN_MARK" "$TRY_MARK" "$GAVEUP_MARK" 2>/dev/null || true
    ;;

  rejoin)
    [ -s "$STATE" ] || exit 0                   # nothing remembered  do nothing
    now=$(date +%s)

    # Grace: remember when the uplink first dropped and hold off, so a DHCP-only
    # blip on the same AP recovers on its own before we force anything.
    if [ ! -f "$DOWN_MARK" ]; then echo "$now" > "$DOWN_MARK"; echo 0 > "$TRY_MARK"; exit 0; fi
    downsince=$(cat "$DOWN_MARK" 2>/dev/null || echo "$now")
    [ $(( now - downsince )) -lt "$GRACE" ] && exit 0

    # Attempt cap: if re-associating hasn't restored the uplink after MAX_TRIES,
    # STOP. Churning a network that won't come up (captive portal, DHCP blocked,
    # out of range) only makes things worse  which is exactly what happened
    # once recovery was aborting. Resets when the link returns (see 'reset').
    tries=$(cat "$TRY_MARK" 2>/dev/null || echo 0)
    if [ "$tries" -ge "$MAX_TRIES" ]; then
      [ -f "$GAVEUP_MARK" ] || { log "gave up after $tries attempts  reconnect Wi-Fi manually"; : > "
     $GAVEUP_MARK"; }
      exit 0
    fi

    # Progressive backoff between attempts: 30s, 60s, 120s, 240s, then cap 300s
    # give DHCP time to actually settle instead of hammering every 30s.
```

```
    backoff=$(( BACKOFF * (2 ** (tries < 4 ? tries : 4)) )); [ "$backoff" -gt 300 ] && backoff=300
    if [ -f "$JOIN_MARK" ]; then
      last=$(cat "$JOIN_MARK" 2>/dev/null || echo 0)
      [ $(( now - last )) -lt "$backoff" ] && exit 0
    fi
    echo "$now" > "$JOIN_MARK"
    echo $(( tries + 1 )) > "$TRY_MARK"

    iface=$(wifi_iface)
    if ! networksetup -getairportpower "$iface" 2>/dev/null | grep -q ': On'; then
      sudo -n /usr/sbin/networksetup -setairportpower "$iface" on >/dev/null 2>&1 || true
      sleep 2
    fi

    log "uplink down  rejoin attempt $((tries + 1))/$MAX_TRIES, remembered network(s) only (iface $iface)"
    while IFS= read -r ssid; do
      [ -z "$ssid" ] && continue
      # -setairportnetwork joins by name using macOS's saved credentials; it
      # only ever targets a network you explicitly remembered. Exit 0 just means
      # the request was accepted  DHCP/uplink recovery is confirmed separately
      # when the watcher sees the interface get an IP and logs RESUMED.
      if sudo -n /usr/sbin/networksetup -setairportnetwork "$iface" "$ssid" >/dev/null 2>&1; then
        log "sent re-associate request for a remembered network (waiting for DHCP)"
        exit 0
      fi
    done < "$STATE"
    log "no remembered network in range (or credentials needed)"
    ;;

  *)
    echo "usage: ${0##*/} {set <SSID>|forget <SSID>|list|rejoin}" >&2
    exit 2
    ;;
esac
```

57 tor-bridges.txt

58 white_list.txt

```
0.client-channel.google.com
1drv.com
2.android.pool.ntp.org
akamaihd.net
akamaitechnologies.com
akamaized.net
amazonaws.com
android.clients.google.com
api-adservices.apple.com
api.ipify.org
api.rlje.net
app-api.ted.com
appleid.apple.com
apps.skype.com
appsbackup-pa.clients6.google.com
appsbackup-pa.googleapis.com
appspot-preview.l.google.com
apt.sonarr.tv
aspnetcdn.com
attestation.xboxlive.com
ax.phobos.apple.com.edgesuite.net
books-analytics-events.apple.com
brightcove.net
c.s-microsoft.com
cdn.cloudflare.net
cdn.embedly.com
cdn.optimizely.com
cdn.vidible.tv
cdn2.optimizely.com
cdn3.optimizely.com
cdnjs.cloudflare.com
cert.mgt.xboxlive.com
```

```
clientconfig.passport.net
clients1.google.com
clients2.google.com
clients3.google.com
clients4.google.com
clients5.google.com
clients6.google.com
connectivitycheck.android.com
connectivitycheck.gstatic.com
continuum.dds.microsoft.com
cpms.spop10.ams.plex.bz
cpms35.spop10.ams.plex.bz
cse.google.com
ctldl.windowsupdate.com
cws.conviva.com
d2c8v52ll5s99u.cloudfront.net
d2gatte9o95jao.cloudfront.net
dashboard.plex.tv
dataplicity.com
def-vef.xboxlive.com
delivery.vidible.tv
dev.virtualearth.net
device.auth.xboxlive.com
display.ugc.bazaarvoice.com
displaycatalog.mp.microsoft.com
dl.delivery.mp.microsoft.com
dl.dropbox.com
dl.dropboxusercontent.com
dns.msftncsi.com
download.sonarr.tv
drift.com
driftt.com
dynupdate.no-ip.com
ecn.dev.virtualearth.net
edge.api.brightcove.com
eds.xboxlive.com
fonts.gstatic.com
forums.sonarr.tv
g.live.com
geo-prod.do.dsp.mp.microsoft.com
geo3.ggpht.com
gfwsl.geforce.com
giphy.com
github.com
github.io
googleapis.com
gravatar.com
gstatic.com
help.ui.xboxlive.com
hls.ted.com
i.ytimg.com
i1.ytimg.com
iadsdk.apple.com
imagesak.secureserver.net
img.vidible.tv
imgix.net
imgs.xkcd.com
instantmessaging-pa.googleapis.com
intercom.io
jquery.com
jsdelivr.net
keystone.mwbsys.com
lastfm-img2.akamaized.net
licensing.xboxlive.com
live.com
livepassdl.conviva.com
login.live.com
login.microsoftonline.com
manifest.googlevideo.com
meta-db-worker02.pop.ric.plex.bz
meta.plex.bz
meta.plex.tv
microsoftonline.com
msftncsi.com
my.plexapp.com
nexusrules.officeapps.live.com
nine.plugins.plexapp.com
no-ip.com
node.plexapp.com
notes-analytics-events.apple.com
notify.xboxlive.com
```

```
npr-news.streaming.adswizz.com
ns1.dropbox.com
ns2.dropbox.com
o1.email.plex.tv
o2.sg0.plex.tv
ocsp.apple.com
office.com
office.net
office365.com
officeclient.microsoft.com
om.cbsi.com
onedrive.live.com
outlook.live.com
outlook.office365.com
pings.conviva.com
placehold.it
placeholdit.imgix.net
players.brightcove.net
pricelist.skype.com
products.office.com
proxy.plex.bz
proxy.plex.tv
proxy02.pop.ord.plex.bz
pubsub.plex.bz
pubsub.plex.tv
raw.githubusercontent.com
redirector.googlevideo.com
res.cloudinary.com
s.gateway.messenger.live.com
s.marketwatch.com
s.youtube.com
s.ytimg.com
s1.wp.com
s2.youtube.com
s3.amazonaws.com
sa.symcb.com
secure.avangate.com
secure.brightcove.com
secure.surveymonkey.com
services.sonarr.tv
skyhook.sonarr.tv
spclient.wg.spotify.com
ssl.p.jwpcdn.com
staging.plex.tv
status.plex.tv
t.co
t0.ssl.ak.dynamic.tiles.virtualearth.net
t0.ssl.ak.tiles.virtualearth.net
tawk.to
tedcdn.com
themoviedb.com
thetvdb.com
tinyurl.com
title.auth.xboxlive.com
title.mgt.xboxlive.com
traffic.libsyn.com
tvdb2.plex.tv
tvthemes.plexapp.com
twimg.com
ui.skype.com
video-stats.l.google.com
videos.vidible.tv
vidtech.cbsinteractive.com
weather-analytics-events.apple.com
widget-cdn.rpxnow.com
win10.ipv6.microsoft.com
wp.com
ws.audioscrobbler.com
www.dataplicity.com
www.googleapis.com
www.msftconnecttest.com
www.msftncsi.com
www.no-ip.com
www.youtube-nocookie.com
xbox.ipv6.microsoft.com
xboxexperiencesprod.experimentation.xboxlive.com
xflight.xboxlive.com
xkms.xboxlive.com
xsts.auth.xboxlive.com
youtu.be
youtube-nocookie.com
```

```
yt3.ggpht.com
zee.cws.conviva.com

# YouTube family  base domains so every subdomain (video stream servers on
# *.googlevideo.com, image hosts on *.ytimg.com, profile thumbnails on
# *.ggpht.com, etc.) is automatically allowlisted via the AdGuard '||domain^'
# wildcard. The specific s.youtube.com / manifest.googlevideo.com entries
# above are kept for clarity but will be optimized away by rule-compiler
# because their parents are now in this list.
youtube.com
googlevideo.com
ytimg.com
ggpht.com
gvt1.com
gvt2.com
gvt3.com
gvt1.net
gvt2.net
gvt3.net

edge-mqtt.facebook.com
gateway.instagram.com
i.instagram.com

facebook.com
fb.com
fbcdn.net
fbsbx.com
messenger.com
whatsapp.com
web.whatsapp.com
instagram.com
cdninstagram.com
static.whatsapp.net
media-bos5-1.cdn.whatsapp.net
media-yyz1-1.cdn.whatsapp.net
webtp.whatsapp.net
mmg.whatsapp.net
g.whatsapp.net
g-fallback.whatsapp.net
```