Nullexit: Unified Project Document

Omar Alaaeldein

July 14, 2026

Contents

1 Project Directory Tree

```
nullexit/
|-- .aiignore
|-- .claude
|   `-- settings.local.json
|-- .gitignore
|-- LICENSE
|-- README.md
|-- Recover-Gateway.applescript
|-- Toggle-Gateway.applescript
|-- agent.md
|-- benchmarks
|   |-- benchmark.sh
|   `-- test_load.py
|-- black_list.txt
|-- cloud-init.yaml
|-- devref.md
|-- diagrams.md
|-- docker
|   `-- tor
|       |-- Dockerfile
|       `-- entrypoint.sh
|-- docker-compose.yml
|-- launchd
|   `-- com.nullexit.wake-recovery.plist
|-- post-rules.txt
|-- recover.sh
|-- scripts
|   |-- Dockerfile.routing-fix
|   |-- Dockerfile.rule-compiler
|   |-- Toggle-Gateway.bat
|   |-- check_circuits.py
|   |-- common.sh
|   |-- crypto.sh
|   |-- diagnose-host-leak.sh
|   |-- dns-proxy.py
|   |-- fix-docker-bridge-collision.sh
|   |-- fix-ssh-delay.sh
|   |-- generate_tex.py
|   |-- logger
|   |   |-- go.mod
|   |   `-- main.go
|   |-- pf.conf
|   |-- reinstall-host-tailscale.sh
|   |-- routing-fix.sh
|   |-- rule-compiler
|   |   |-- go.mod
|   |   `-- main.go
|   |-- setup-common.sh
|   |-- setup-linux.sh
|   |-- unlock-files.sh
|   `-- watcher.sh
|-- setup.sh
```

```
|-- toggle.sh
|-- tor-bridges.txt
`-- white_list.txt
```

2 README.md

```
# nullexit: Tailscale + Cloudflare WARP Docker Gateway

> **Last updated:** July 12, 2026  [Architecture & Flow Diagrams ](./diagrams.md)  [Full Dev Reference ](./devref.md)

**nullexit** is a chained network gateway that routes all Tailscale exit-node traffic through a Cloudflare WARP VPN tunnel  double-encrypting every packet, hiding your ISP metadata, and providing network-wide DNS ad-blocking (AdGuard Home) and kernel-level IP threat blocking (`ipset`/`iptables`) for every device on your mesh.

---

## 1. Prerequisites
- Docker and Docker Compose installed.
- A Tailscale account.
- **macOS:** Install the standalone Tailscale CLI via Homebrew (`brew install tailscale`) and register `tailscaled` as a system service (`brew services start tailscale`). No `.app` GUI install required.
- **OS & Python:** Developed and tested against **macOS 26.5.2**. The project has zero external Python dependencies (no `requirements.txt`) and explicitly targets the default Apple-provided **Python 3.9.6** (`/usr/bin/python3`). While the scripts could function on newer versions, `nullexit` explicitly avoids experimenting with or modifying the built-in system Python environment to prevent unexpected OS-level side effects.

## 2. Installation & Setup

Run the interactive setup script for your OS. It downloads `wgcf`, generates fresh Cloudflare WARP WireGuard keys, prompts for a Tailscale Auth Key, and writes your `.env`.

```bash
./setup.sh          # macOS
./scripts/setup-linux.sh  # Linux
```

Reference `.env`
This is a quick reference of the common variables. For the **complete, authoritative list** of every `.env` variable (including all optional/advanced settings), see `devref.md 7`.
```env
WIREGUARD_PRIVATE_KEY=<Generated>
WIREGUARD_PUBLIC_KEY=<Generated>
WIREGUARD_ADDRESSES=<Generated>
TS_AUTHKEY=tskey-auth-...
GATEWAY_RULE_PROFILE=medium        # light | medium | heavy
# Safe MSS for double-tunneled traffic (Tailscale + WARP). Change to 1180 if you experience slow speeds and have a healthy path.
GATEWAY_MSS=1120
# Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own internet is NOT hijacked).
GATEWAY_HIJACK_HOST=true
GATEWAY_USE_EXIT_NODE=true
WARP_FAIL_THRESHOLD=6              # consecutive warp=off polls before auto-shutdown (default 6 = 30s)
KILL_SWITCH=false                  # enforce strict PF lock that breaks SSH if VPN fails
# On campus/enterprise networks (WPA2-Enterprise / 802.1x), AP Client Isolation blocks direct
# LAN traffic between devices. Tailscale attempts a local P2P upgrade anyway, which causes
# macOS gvproxy to SNAT-mangle the reply's source port  poisoning the phone's WireGuard
# endpoint and blackholing 100% of its traffic. Setting this to false drops those local
# hole-punch packets so the phone safely falls back to Tailscale DERP relays.
# watcher.sh auto-detects 802.1x and AP isolation on every Wi-Fi change and overrides
# this setting automatically  you only need to set this manually if auto-detection fails.
TAILSCALE_ALLOW_LAN_P2P=false     # auto-managed; true only on trusted hotspots/home networks
ADGUARD_USER=admin                 # Username for AdGuard Home (defaults to admin if not set)
ADGUARD_PASSWORD=nullexit          # Shared password for AdGuard Home and Tor ControlPort (defaults to nullexit if not set)
BLOCKED_COUNTRIES="kp il"          # dynamically block country IP ranges via ipdeny.com
TOR_EXCLUDE_EXIT_NODES=""          # comma-separated country codes to exclude as Tor exit nodes (e.g. {us},{gb})
DEBUG_TRACE=false                  # true = mirror a full command-by-command trace of the toggle to output.log (see 9)
```

3. Deploy the Gateway

macOS (Recommended):
```bash
./toggle.sh
```
Running it again stops the gateway and restores your network. Handles the full lifecycle: Colima VM, containers, DNS hijacking, Tailscale exit-node routing, rule compilation, and sleep prevention.

```

```
Linux / Manual:
```bash
docker compose up -d
```

4. Post-Deployment

Approve the Exit Node
1. Go to [Tailscale Admin Machines](https://login.tailscale.com/admin/machines).
2. Find the `tailscale` device `...` **Edit route settings** enable as Exit Node.

Verify Traffic Flow
On any Tailscale client device, select this gateway as the exit node, then visit api.ipify.org or run:
```bash
curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
# expect: warp=on
```

AdGuard DNS Setup

To ensure client devices correctly route DNS through the AdGuard container (and do not bypass it via `tailscaled`'s internal resolver), you **must** configure your Tailscale Admin Console:
1. Go to the **DNS** tab in the [Tailscale Admin Console](https://login.tailscale.com/admin/dns).
2. Enable **MagicDNS**.
3. Under **Nameservers**, add a **Custom** nameserver and enter the gateway's Tailscale IPv4 address (run `tailscale ip -4` on the gateway machine to find it, e.g., `100.x.x.x`).
4. Click the **three dots (...)** next to the nameserver you just added and enable **"Use with exit node"**.
5. Delete any other Global Nameservers (like Google or Cloudflare).
6. Toggle ON **Override local DNS**.

Traffic will now be correctly forced into the AdGuard sinkhole, and mobile devices will be blocked from bypassing it via encrypted DNS (DoH/DoT).

> [!WARNING]
> **Disable "Private DNS" on mobile devices.** Android's Private DNS (DoT/DoH) bypasses port 53 interception entirely. Go to `Settings  Connections  More connection settings  Private DNS  Off`.

> [!WARNING]
> **Bandwidth doubling:** Streaming 1GB of video on a remote device consumes 1GB download + 1GB upload on the host's network. Be mindful on metered connections.

AdGuard Dashboard
Navigate to `http://<gateway-tailscale-ip>:3000` credentials: **`admin` / `nullexit`**.

Access Any Mesh Device From Anywhere (SFTP/SSH)

Once the gateway exit node is online, **any device on your Tailscale mesh can reach any other mesh device, anywhere in the world** as long as both devices are on the mesh and the target device has an SSH server running. No port forwarding, no public IPs, no firewall holes. You can SSH into any device using its MagicDNS hostname (e.g. `ssh username@my-macbook`) or its assigned Tailscale IP (`100.x.x.x`).

> **Security Tip (Zero Local Attack Surface):** It is highly recommended to disable macOS's native "Remote Login" and "Screen Sharing" in System Settings to prevent exposing those ports (22, 5900) to physical Wi-Fi networks (which scanners can fingerprint). The gateway enforces `tailscale up --ssh=true` and the `pf` Kill-Switch automatically blocks local Wi-Fi access to these ports. Since both MagicDNS and Tailscale IPs route through the exact same encrypted tunnel, they are equally securebut **MagicDNS is recommended simply for ease of use.**

Extreme OPSEC: The Secret Tor-over-WARP Proxy
`nullexit` includes a completely isolated, headless Tor infrastructure proxy designed for terminals and hacking tools (e.g., `curl`, `sqlmap`, `nmap`).
To protect against local malware fingerprinting:
1. **Procedural Ports:** The proxy ports are randomized into the ephemeral range (10000-65000) on every single boot using a strict kernel socket collision-avoidance check (`netstat`).
2. **Honey-Port Tripwire:** The gateway spawns a fake trap port. If any malware attempts a sequential port-scan on `localhost` to find the proxy, it trips the Honey-Port, instantly triggering a macOS desktop notification and a critical log alert.

> [!NOTE]
> **Threat Model limitation:** Port randomization and the Honey-Port only defeat blind, generic port scanners. They do not protect against targeted local malware that runs `lsof -i` or reads the `/tmp/nullexit-ports.env` file directly. To actually prevent malicious local processes from reaching the proxy, you must run a host-level outbound firewall (like LuLu 4.3.0+ or Little Snitch) with **localhost/loopback filtering explicitly enabled** to gate access by process identity.

To use the Tor proxy, run `cat /tmp/nullexit-ports.env` to find your `$TOR_SOCKS_PORT` for the current session, and point your hardened browser (LibreWolf/Tor Browser) or CLI tool to `127.0.0.1:$TOR_SOCKS_PORT`.

> [!WARNING]
> **Preventing DNS Leaks (SOCKS5 vs. SOCKS5-Hostname):**
```

```
> When routing command-line tools through Tor, you **must** force the tool to delegate DNS resolution to the Tor proxy. Using plain SOCKS5 (e.g., `socks5://` or curl's `--socks5` flag) will resolve hostnames *locally* using the host's DNS settings (AdGuard/WARP) before establishing the connection. While this DNS traffic is encrypted inside the WARP tunnel, the destination domain name is still leaked to AdGuard and Cloudflare, undermining your anonymity.
>
> Use these correct protocols and configurations:
> - **`curl`:** Use `--socks5-hostname` or the `socks5h://` scheme (e.g., `curl -x socks5h://127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org`). Do **not** use `--socks5` or `socks5://`.
> - **`sqlmap`:** Pass the proxy URL using the `socks5h://` scheme to ensure remote resolution: `--proxy="socks5h://127.0.0.1:$TOR_SOCKS_PORT"`.
> - **`nmap`:** Nmap's built-in `--proxies` option resolves hostnames locally. To safely scan targets without DNS leaks, tunnel it using `proxychains-ng` configured with `proxy_dns` enabled (add `socks5 127.0.0.1 <PORT>` to `/etc/proxychains.conf` or a local config, then run `proxychains4 nmap <args>`).


#### Transparent .onion Browsing & Dark Web Access
In addition to the randomized SOCKS5 proxy port, `nullexit` provides seamless, transparent `.onion` domain browsing for **all devices** on your Tailscale mesh:
- **Zero-Config DNS:** AdGuard Home automatically intercepts queries for `.onion` domains and resolves them to Tor's virtual mapping subnet.
- **Tailscale Auto-Routing:** The virtual subnet is mapped to the `198.18.0.0/15` benchmark range (RFC 2544). Since this range is not a private IP subnet, Tailscale's Exit Node mode automatically routes it to the gateway (bypassing local LAN exclusions).
- **Transparent NAT Proxying:** Kernel firewall rules in the gateway redirect all TCP traffic destined for `198.18.0.0/15` directly into Tor's transparent port (`9040`).
- **Exit Node Exclusion:** Supports the ability to exclude specific countries from being used as Tor exit nodes via the `TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`.

##### Enable in Web Browsers
To prevent accidental DNS leaks, web browsers block `.onion` lookups by default. To browse dark web sites natively:
- **Firefox:** Go to `about:config`, search for `network.dns.blockDotOnion`, and set it to `false`. Type any onion address (e.g. `duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion`) directly into the URL bar and it will load.
- **Mobile Browsers (iOS/Android):** Connect to your Tailscale exit node, open Firefox (with blockDotOnion disabled), and browse onion sites natively on the go.

#### Hardening: Using obfs4 Tor Bridges
If you want to disguise your Tor traffic from the WARP exit node provider, you can optionally enable Tor bridges:
1. Obtain private `obfs4` bridge lines from [bridges.torproject.org](https://bridges.torproject.org) or the official Telegram bot.
2. Create a file named `tor-bridges.txt` in the root of the `nullexit` folder and paste your bridge lines inside it (one per line).
3. Set `TOR_USE_BRIDGES=true` in your `.env` file and restart the gateway. The proxy will refuse to start if the bridges file is empty or missing.

#### Example: Browse your Android phone's files from your Mac


1. **On your Android phone**, install [Termux](https://termux.dev/) and [Termux:Boot](https://f-droid.org/packages/com.termux.boot/) from F-Droid.
2. In Termux, install and start the SSH server:
   ```bash
   pkg install openssh
   sshd
   ```
   Termux defaults to **port 8022** (ports below 1024 require root on Android).
3. Find your phone's Tailscale IP:
   ```bash
   tailscale ip -4   # e.g. 100.87.42.87
   ```
4. **Stop Android from killing Termux** do ALL of these (Samsung One UI is especially aggressive):
   - **Settings Apps Termux Battery Unrestricted**
   - **Device Care Battery Background usage limits Sleeping apps / Deep sleeping apps** remove Termux if listed
   - In Termux, run `termux-wake-lock` to hold a persistent wakelock
5. **On your Mac/PC**, open [Cyberduck](https://cyberduck.io/) (or any SFTP client: WinSCP, FileZilla, FE File Explorer on iOS) and connect to:
   - **Server:** `100.87.42.87` (your phone's Tailscale IP)
   - **Port:** `8022`
   - **Protocol:** SFTP (SSH File Transfer Protocol)
   - **Username:** (your Termux username, usually `u0_a123` run `whoami` in Termux to check)
   - **Password:** (your Termux password set with `passwd` in Termux if you haven't already)

That's it you can now browse, download, and upload files to your phone from any device on the mesh, no matter where either device is physically located. This works identically for any mesh device: Mac Mac, Windows Linux, phone NAS, etc.

```

```
> **Troubleshooting:** If the connection hangs for ~30 seconds before failing, your phone's SSH server isn't running. Open Termux and run `sshd` again. If Android keeps killing it despite Unrestricted battery, the `termux-wake-lock` command is your fix.

---

5. Multi-Layer Firewall & Kill-Switch

Layer 1 DNS Sinkhole (AdGuard Home)
Blocks ads and tracking domains before they are downloaded. In automated Lighthouse tests on ad-heavy sites:
- **Time to Interactive:** ~22% faster (51.0s 41.8s)
- **Speed Index:** ~20% faster (15.3s 12.8s)

Customize blocking via `black_list.txt` and `white_list.txt`. Rule profiles (`light` / `medium` / `heavy`) are set in `.env`.

Layer 2 Kernel IP Firewall (`ipset`/`iptables`)
On every startup, the `rule-compiler` fetches threat-intelligence feeds (Spamhaus DROP, Feodo Tracker, Emerging Threats, CINS) and compiles **~16,700 unique malicious IPs/CIDRs** into the kernel `FORWARD` chain blocking both outbound C2 connections and inbound attack traffic. Zero configuration required.

Layer 3 Geo-IP Blocking
Dynamically blocks all traffic to and from specific countries using live IP ranges from `ipdeny.com`.
To add countries to your blocklist, open your `.env` file and set the `BLOCKED_COUNTRIES` variable with 2-letter ISO country codes (e.g. `BLOCKED_COUNTRIES="kp il cn ru"`), then restart the gateway.

Layer 4 macOS PF Kill-Switch
A native macOS Packet Filter (`pf`) kill-switch that enforces a strict default-deny policy on your physical Wi-Fi interface (`en0`). When `KILL_SWITCH=true` is set in your `.env`, it drops all outgoing traffic except the Cloudflare WARP endpoints and Tailscale DERP relays.

- **Failsafe Design:** If the VPN tunnel crashes or Docker dies, your Mac completely loses internet access rather than leaking your real IP.
- **Remote-Access Safe:** Carefully designed to whitelist Tailscale's control plane (`192.200.0.0/24`), `utun*` tunnel traffic, and local LAN subnets. This guarantees that your remote SSH sessions and local AirDrop will survive even when the kill switch engages.
- **Setup Requirement:** The toggle scripts need permission to manipulate the network and firewall in the background without prompting you for a password. You must configure your passwordless sudo config by running exactly:
  ```bash
  echo "$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /usr/bin/dscacheutil, /usr/bin/killall, /usr/bin/pkill, /bin/kill, /sbin/route, /sbin/ifconfig, /usr/bin/true, /opt/homebrew/bin/brew, /usr/local/bin/brew, /usr/bin/python3 -I $PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I $PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I $PWD/scripts/dns-proxy.py" | sudo tee /etc/sudoers.d/nullexit
  ```

---

6. Toggle Scripts

macOS `toggle.sh`
Fully automated lifecycle script. Also supports a native `.app` launcher:
```bash
osacompile -o "Toggle Gateway.app" Toggle-Gateway.applescript
```

Optional `.env` settings (common subset; the full canonical table lives in `devref.md 7`):
- `GATEWAY_BYPASS_PING=true` proceed even if pre-flight connectivity checks fail.
- `GATEWAY_USE_EXIT_NODE=false` DNS-only mode, skip exit-node routing.
- `GATEWAY_HIJACK_HOST=false` skip DNS hijacking on the host (provides VPN/adblocking only to external Tailscale peers).
- `GATEWAY_MSS=1180` override the TCP MSS clamp for the double-tunnel (default `1120`; raise to `1180` for more speed on a healthy path).
- `HOST_LEAK_PROBE=false` disable background host-egress leak prober (default: true).
- `HOST_MTU=1200` override the host Tailscale interface MTU. Defaults to 1200 to fit inside the 1280-byte WARP tunnel.
- `KILL_SWITCH=true` enforce strict network lock that breaks SSH if the VPN fails.
- `STOP_COLIMA_ON_EXIT=true` fully shut down the Colima VM on toggle-off (saves battery on dedicated hosts).
- `TAILSCALE_ALLOW_LAN_P2P=true` allow direct Tailscale LAN P2P connections between devices on the same network.
  - **Default: `false`.** On campus or enterprise Wi-Fi (WPA2-Enterprise / 802.1x), AP Client Isolation blocks local device-to-device traffic. When Tailscale still attempts a local P2P upgrade, macOS's `gvproxy` layer SNAT-mangles the reply's source port. WireGuard's endpoint roaming treats this as a legitimate address change and updates the phone's outbound path to the newly-mangled port which has no listener. This blackholes 100% of the phone's traffic while the DERP relay path is abandoned. Setting this to `false` preemptively drops all RFC1918-destined Tailscale UDP from the gateway so the phone receives a clean failure and safely falls back to DERP.
  - **Auto-managed:** `watcher.sh` automatically detects WPA2-Enterprise via `system_profiler SPAirPortDataType` (the `airport` binary was removed in macOS 14.4) and probes for AP isolation by pinging the default gateway (`route -n get 1.1.1.1`) on every network change. It writes the detected value to `.lan_p2p_detected` in the repo root, which `routing-fix.sh` re-reads every 30 seconds **no restart required
```

```
     .** Only set this manually if auto-detection fails. See `devref.md 15.7.1` for the full SNAT endpoint
     poisoning analysis and `devref.md 15.11.6` for the `airport` removal background. *(Note: Direct P2P between
      the gateway container and the host Mac itself is always permitted via the Docker bridge to bypass DERP
     relays, regardless of this setting).*
  - Set to `true` on trusted home networks or Windows/mobile hotspots (no AP isolation) for a faster, lower-
     latency direct path.
- `WARP_FAIL_THRESHOLD=3`  number of consecutive failed checks before forcing a gateway teardown (default: 6
     checks = 30s).
- `WARP_ENDPOINT_1` / `WARP_ENDPOINT_2`  override the Cloudflare WARP WireGuard endpoints.

### Linux  `toggle.sh` / `recover.sh`
```bash
./toggle.sh    # toggle ON/OFF   (cross-platform, dispatches on $OSTYPE)
./recover.sh   # recovery tool   (cross-platform)
```

Sleep Prevention (macOS)
When the gateway is ON, `caffeinate -i` runs in the background to prevent idle system sleep, keeping Docker and
      all services alive even on battery. The display can still sleep normally.

Auto-Recovery (macOS post-wake / post-roam)
Install the launchd LaunchAgent so the gateway self-heals after sleep/wake and Wi-Fi roams:
```bash
cp ./launchd/com.nullexit.wake-recovery.plist ~/Library/LaunchAgents/
sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```
See `devref.md 15.5.1` for the full deep dive.

---

7. Networking Notes

- **Port conflicts:** None. Tailscale uses dynamic UDP ports; WARP uses outbound UDP:2408. The gateway binds
     `0.0.0.0:41642/udp` on the host for Tailscale's WireGuard listener this is intentional and required to
     receive inbound hole-punch packets from peers on the internet. All other internal ports (AdGuard :5335,
     SOCKS5 :1080) are bound to `127.0.0.1` only and are not reachable from the LAN.
- **Tailscale P2P vs DERP:** On the same Wi-Fi, Tailscale establishes a fast P2P connection. On cellular (CGNAT
     ), it always falls back to a DERP relay (~80-200ms). See `devref.md 5.1`.
- **AirDrop / AirPlay:** Unaffected the gateway only touches standard Wi-Fi/Ethernet interfaces, not `awdl0`.
- **VPN coexistence:** Do **not** run a local VPN client (WARP, Mullvad, NordVPN) simultaneously with the exit
     node. Both fight for the default route and will blackhole your connection.
- **Banking & financial sites:** May intermittently block logins through WARP. Banks use anti-fraud databases
     that flag datacenter IP ranges (like Cloudflare's `104.28.x.x`) as VPN/proxy traffic. Success fluctuates
     depending on which WARP IP you land on. If a site blocks you, either disable the exit node temporarily for
     that session or use the bank's mobile app (which often uses different fraud detection).
- **MSS clamping (Bufferbloat):** Double-tunnelling adds ~120-160 bytes of overhead per packet. This is now
     automatically handled natively via the macOS `pf` firewall, which universally applies TCP MSS Clamping to
     all outbound packets to strictly prevent MTU fragmentation and bufferbloat. No manual `.env` configuration
     is required. Additionally, the host's Tailscale interface MTU is explicitly lowered to 1200 (configurable
     via `HOST_MTU` in `.env`) to guarantee packets fit inside the container's WARP tunnel without fragmenting.

---

8. Privacy & Security Hardening

The stack shifts trust across four layers: WARP hides traffic from your ISP, AdGuard blocks tracking at DNS,
     Tailscale encrypts device-to-device, and `ipset` blocks known malicious IPs at the kernel.

For higher security, the gateway supports these drop-in upgrades:

Layer	Default	Hardened
VPN Exit	Cloudflare WARP (US)	Mullvad (Swedish, no-logs, anonymous payment)
DNS Resolver	AdGuard via WARP	AdGuard via Mullvad DoH
Coordination	Tailscale (US)	Headscale (self-hosted)
Auth	Persistent OAuth keys	Ephemeral auth keys
Content	WireGuard asymmetric	WireGuard + Pre-Shared Keys (PSK)

Python Runtime Airgap (Isolated Mode)
Never install third-party `pip` packages (like `requests` or `pyyaml`) into the Python environment that `
     nullexit` uses. All python scripts in this project (e.g., the `dns-proxy.py` daemon) are strictly
     engineered to run on the zero-dependency Python Standard Library. To strictly isolate the environment from
     supply-chain attacks, `nullexit` executes these scripts using **Python's Isolated Mode** (`python3 -I`).
     This deliberately blinds the execution environment to any malicious user-installed `pip` packages on the
     host, permanently air-gapping the gateway from PyPI.

See `devref.md 9` for the full threat model, Snowden programme analysis, and trust assumptions.

```

```
## 9. Debugging

- **`output.log`**  All stderr from `toggle.sh`, `recover.sh`, `setup.sh`, and the rule-compiler is written here. The WARP Watcher (`WARP_FAIL_THRESHOLD`) and the Host Leak Probe (`HOST_LEAK_PROBE`) also write all their events here  `LEAK`, `ROTATE`, `HOST-PROBE failed`, and WARP shutdown notices. Check this first.
  - Every run now records lifecycle breadcrumbs regardless of `DEBUG_TRACE`: an `EXEC:` / `EXIT <code>:` line brackets each heavyweight command (`colima start`, `docker compose up -d --build`, `docker compose down`), and an `EXIT: code=<N> phase=<init|stop|start>` line is written when the script exits **for any reason** including a killed terminal, `pkill`, or OOM. This closes a prior blind spot where an interrupted **START** (e.g. a `--restart` race while Wi-Fi was re-associating) left the log ending abruptly after the teardown with no indication of what failed. If a toggle "does nothing" or the gateway ends up half-up, `grep -E ' EXEC:|EXIT ' output.log | tail` shows the last command run and its result.
- **`DEBUG_TRACE=true`** (in `.env`)  Full command-by-command execution trace of `toggle.sh`, mirrored to `output.log` for deep post-mortems of intermittent failures. Each traced line is timestamped and tagged with `file:line:function` (via a custom `PS4`), so you can follow the *exact* code path a run took  including subshells, Docker/Colima calls, and where it aborted. Off by default because it is verbose (the log grows quickly); enable it only while reproducing a specific problem, then set it back to `false`. Implementation note: it prefers bash's `BASH_XTRACEFD` (bash  4.1, e.g. most Linux) to send the trace to the log while keeping your terminal clean; on macOS's stock `/bin/bash` 3.2 it transparently falls back to teeing stderr to the log (trace also appears in the terminal). It never uses the dynamic-fd (`exec {var}>>`) form, so it is safe on 3.2. To read a trace: `grep -nE '^\+ ' output.log`.
- **`bash scripts/diagnose-host-leak.sh`**  One-shot host-routing diagnostic. Runs 8 checks and classifies your state into one of four known scenarios (System Extension conflict, SOCKS5 fallback, IPv6 leak, route-table freeze) or OK, and prints the exact fix command. Check 5 uses `route -n get 1.1.1.1` to ask the kernel which interface real traffic actually resolves through (more accurate than `netstat -rn | grep default` on macOS). Pass `--fix` to apply the remediation automatically. Pass `--watch` (or `--watch 30`) to run a full baseline diagnostic then continuously monitor warp/IPv6/default-route for leaks, alerting on any state change.
- **`bash scripts/host-leak-probe.sh`**  Sub-second host-egress prober (enabled automatically via `HOST_LEAK_PROBE=true` in `.env`). Polls `cdn-cgi/trace` every 300ms directly from the host NIC  not via `docker exec`  so it catches flash-leaks invisible to the in-container WARP Watcher. All state changes, curl errors, and probe timeouts land in `output.log`. Grep with `grep LEAK output.log`.
- **`bash scripts/fix-docker-bridge-collision.sh`**  One-shot fix for Docker bridge IP conflicts. If your local Wi-Fi router assigns you an IP in the `172.17.0.0/16` range, it will violently collide with Docker's default internal bridge, permanently killing your internet. This script detects the collision and auto-patches your Colima VM to use a safe subnet (e.g. `10.200.0.1/24`).
- **`devref.md`**  Complete architecture reference, routing deep-dives, and full resolved-issues log. Read this before modifying any routing or firewall logic.

> [!CAUTION]
> **Two infinite routing loops are possible.** (1) If a hotspot device providing internet to the gateway host is *also* using the gateway as its exit node, packets bounce forever between the two. (2) When the host's default route is redirected to `utun*`, WARP endpoint packets can loop back into the tunnel. Both are documented in `devref.md 15.2.1` and `15.2.2` and are automatically handled by `toggle.sh`.

---

## 10. Unified Project Document (Quine)

The project includes a `generate_tex.py` script that compiles the entire source code, configurations, and documentation into a single, unified LaTeX document (`nullexit_unified.tex` / `.pdf`).

Funnily enough, this document acts as a "project-level quine." It encapsulates the entirety of its own source code, its documentation, and the exact Python code required to generate itself. It serves as a self-contained, long-term archiving strategyproviding everything needed to reconstruct and understand the `nullexit` system from a single file.

---

## 11. Acknowledgements
- **[SyameimaruKoa](https://github.com/SyameimaruKoa):** Dual-stack TCP MSS clamping, `SIGHUP` state-tracking in the routing sidecar, and `TS_AUTH_ONCE` integration.

## 12. License
GNU Affero General Public License v3. See [LICENSE](./LICENSE).

## 13. Changelog

- **July 12, 2026**  Tor ControlPort dynamic password authentication (unified with `ADGUARD_PASSWORD`), enabling live circuit inspection via `/sweep`. See `devref.md 15.10.2`.
- **July 11, 2026**  Tor container hardening: `obfs4proxy` restored (base image  `debian:bookworm-slim`), robust bridge-file validation, and image pinning. See `devref.md 15.10.1`.
- **July 10, 2026**  Roaming/startup stabilization suite: fixed the Wi-Fi-roam host-IP leak (raw ISP `132.x`) caused by the WARP Watcher racing post-wake recovery, plus a batch of pre-flight/concurrency/kill-switch/Tailscale-flag bugs and the macOS SNAT endpoint-poisoning P2P blackhole. Full post-mortems in the Incident Log  see `devref.md 15.2.7`, `15.5.3``15.5.5`, `15.2.8``15.2.9`, `15.7.1`, `15.7.3``15.7.5`, `15.12.18`.
- **July 6, 2026**  Edge-case security and stability hardening: DNS Watcher interface independence, robust lock-file PID verification, fail-closed crypto integrity check, strict `iptables` pinning for tailscale0tun0, Docker port binding to `127.0.0.1` to prevent LAN exposure, routing verification via `netstat -nr`. See `devref.md 17`.
```

```
- **July 4, 2026**  Colima bridged networking, SOCKS5 proxy migrated natively to Tailscale, headless prompt crashes fixed, proxy bypass domains added to fix LAN P2P. Python dependencies completely eliminated: `logger` and `rule-compiler` ported to blazing-fast Go multi-stage Docker builds. Added `crypto.sh` to enforce cryptographic signature validation on all core bash scripts. Massively refactored `toggle.sh` and `recover.sh` to eliminate ~200 lines of duplicated code by migrating network and process logic to `scripts/common.sh`. Added a concurrency mutex to `toggle.sh`, resolved unbound `TS_AUTHKEY` check crashes on setup, and integrated Go validation to the CI pipeline.
- **July 3, 2026**  **Critical Security Updates:** Addressed the overnight silent IP leak and improved geo-blocking. Introduced `scripts/host-leak-probe.sh` (a sub-second host-egress leak detector) and a statistical threshold auto-shutdown watcher. Added `unlock-files.sh` to safely reset file permissions. Resolved numerous startup regressions and system bugs.
- **July 2, 2026**  Two infinite routing loops resolved (`devref.md 15.2.2`): host-VM tunnel loop (WARP bypass routes now via gateway IP + manual `utun*` redirection) and Docker Compose subnet takeover (`10.200.1.0/24` lock). `--accept-routes=true` now explicit on all `tailscale up --reset` calls.
- **July 2, 2026**  Auto-recovery daemon documented in 6 above (`scripts/watcher.sh` + launchd plist). See `devref.md 15.5.1`.
- **July 2, 2026**  Linux scripts (`scripts/toggle-linux.sh`, `recover-linux.sh`, `setup-linux.sh`) documented in 6.
- **July 1, 2026**  `recover.sh --post-wake` ships; wired to watcher daemon. Triggered on every sleep/wake or network-state change.
- **June 28, 2026**  8 internal scripts moved to `scripts/`. User-facing files (`toggle.sh`, `recover.sh`, `setup.sh`, `docker-compose.yml`) stay at repo root.
- **June 28, 2026**  All hardcoded install paths removed. `SCRIPT_DIR`-relative resolution and `__NULLEXIT_HOME__` placeholder used throughout.

### Cryptographic Integrity Verification

nullexit uses a built-in cryptographic integrity checker designed to provide tamper-*evidence* against accidental edits, logic bugs, or non-privileged malware. During setup, a 256-bit `NULLEXIT_SEED` is injected into your `.env` file. The core entrypoint scripts (`toggle.sh`, `recover.sh`, etc.) are hashed using HMAC-SHA256, and their signatures are stored in `.signatures`.

*(Note: This is not a strict boundary against an attacker who already has write access to the repo, as they could simply read the seed and re-sign the scripts. It is designed as a strict footgun-catcher).*

If any script is modified without authorization, the gateway will loudly fail to boot. If you make intentional edits to the bash scripts, you must re-sign them by running:
```bash
./scripts/crypto.sh --sign
```
```

3 agent.md

```
# nullexit  Detailed Agent Analysis

> Complete per-file breakdown of every function, constant, duplication, and bug found.
> **Note (July 2026):** `sync-rules.py` and `logger.py` have been fully ported to Go (`scripts/rule-compiler/main.go` and `scripts/logger/main.go`) to dramatically improve performance and reduce container footprint via multi-stage Alpine builds. The analysis below reflects their previous Python states.

##  Important Agent Instruction
- **NEVER use `cat << EOF` or terminal redirections to write or modify files.** You must always use your native file-editing tools (`replace_file_content` or `multi_replace_file_content`). Using `cat EOF` leads to duplicated text, broken markdown headers, and structural corruption (which previously destroyed the section numbering in `devref.md`).
- **Always run `bash scripts/crypto.sh --sign` after making any modifications to the core bash scripts.** This is required to update the cryptographic HMAC-SHA256 signatures; otherwise, the startup integrity checks will fail and block execution.
- **Always run `git diff` and print the changes to the user *before* requesting or executing any git commit command.** The user must be shown the diff so they can review and understand the changes. Based on this diff, formulate a highly precise, detailed, and descriptive commit message (explaining exactly what was changed and why) rather than a generic summary, and present it to the user alongside the diff.
- **Always run `git status` before `git diff` to identify all modified and untracked files, ensuring no files (such as new scripts or configuration assets) are missed before formulating commit diffs and staging.**
- **Always use the `--restart` flag when asked to restart the toggle (e.g., run `./toggle.sh --restart`).**
- **NEVER execute any `sudo` commands (especially `sudo route`, `sudo dscacheutil`, or any network-modifying commands) without explicitly explaining to the user what the command does and why it is needed FIRST. Do not assume implicit permission. Wait for the user's approval before running it.**
- **When fixing an error and using logs to debug, always show the user the reasoning and the specific logs that support this theory.**
- **When the user tells you to "push", this means read git diffs and prepare commit messages that correspond to these changes (do not ignore any changes). If the changes can be atomized into multiple commits for easier understanding, do so.**
- **When the user says "latex this", it means you should run `python3 scripts/generate_tex.py` to regenerate the unified LaTeX document.**
- **NEVER apply changes to running gateway containers directly via `docker compose up` or `docker compose restart`.** Recreating or restarting containers (especially the `warp` container which hosts the network
```

```
     namespace) breaks the shared network stack for all other services (`adguardhome`, `tailscale`, `routing-fix
     `), severing host connectivity and freezing macOS DNS. If you need to apply changes or rebuild containers,
     cleanly stop the gateway first using `./toggle.sh`, or trigger `recover.sh --post-wake` to rebuild and re-
     verify the network stack safely in the correct dependency order. Do NOT run `./toggle.sh --restart` or any
     network-rebuilding commands yourself if the execution could sever host network access; instead, explain the
      changes and ask the USER to run `./toggle.sh --restart` (or `./toggle.sh` stop and start) themselves to
     safely apply the modifications.

##  How to Verify Gateway is Working
Whenever modifications are made to the gateway scripts, routing, or containers, execute these verification
     checks in order:
1. **Container Health:** Run `docker compose ps` to ensure all containers (`warp`, `adguardhome`, `routing-fix
     `, `tailscale`) are `running` and `healthy`.
2. **DNS Hijacking Status:** Run `networksetup -getdnsservers "Wi-Fi"` (or appropriate network service) and
     ensure it is set exclusively to the gateway's Tailscale static IP (typically `100.100.21.8`).
3. **DNS Query Interception:** Run `dig @100.100.21.8 google.com` (using the gateway static IP from `.
     gateway_ip`) to ensure AdGuard Home is active, intercepting, and resolving queries.
4. **Double-Tunnel Egress:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp` to verify the
     egress is routed through Cloudflare WARP (`warp=on`).
5. **Host Leak Monitoring:** Run `tail -n 30 output.log` and verify the background watchers and host leak
     prober are reporting healthy status without `LEAK` warnings.

---

## Part 1: macOS Main Scripts

### toggle.sh (1281 lines, 56KB)

**Functions (27 total):**

| # | Function | Lines | Purpose |
|---|----------|-------|---------|
| 1 | `run_with_timeout()` | L3064 | Run a command with a timeout watchdog; kills after N seconds |
| 2 | `run_gui_cmd()` | L6880 | Run a GUI command as the logged-in console user |
| 3 | `write_gateway_active_marker()` | L8890 | Write timestamp to `/tmp/nullexit-gateway-active.marker` |
| 4 | `clear_gateway_active_marker()` | L9295 | Remove gateway active marker + TUNNEL_FAILED_CLOSED.marker |
| 5 | `start_sleep_prevention()` | L108154 | Start `caffeinate` in background with shutdown trap to revert DNS |
| 6 | `stop_sleep_prevention()` | L157167 | Stop caffeinate process via PID file |
| 7 | `start_dns_watcher()` | L173191 | Background loop to re-hijack DNS every 30s for Wi-Fi roaming |
| 8 | `stop_dns_watcher()` | L193203 | Stop DNS watcher via PID file |
| 9 | `start_warp_watcher()` | L220279 | Background WARP liveness monitor; triggers recover.sh after N consecutive failures |
| 10 | `stop_warp_watcher()` | L281291 | Stop WARP watcher via PID file |
| 11 | `start_host_leak_probe()` | L297310 | Start host-egress leak prober (300ms polling) |
| 12 | `stop_host_leak_probe()` | L312322 | Stop host leak probe via PID file |
| 13 | `cleanup_handler()` | L325382 | Error/signal handler: resets DNS, stops all watchers, tears down containers |
| 14 | `restart_tailscaled_daemon()` | L449463 | Restart tailscaled via brew services (user then system fallback) |
| 15 | `disconnect_tailscale_host()` | L465477 | Disconnect host Tailscale with retry on failure |
| 16 | `get_active_service()` | L480508 | Detect active macOS network service name (e.g., "Wi-Fi") |
| 17 | `add_warp_bypass_routes()` | L524546 | Add static routes for WARP endpoints (162.159.192.1, 162.159.193.1) |
| 18 | `remove_warp_bypass_routes()` | L548552 | Remove WARP bypass routes |
| 19 | `setup_exit_node_routing()` | L558577 | Override default route to point through Tailscale utun interface |
| 20 | `reset_dns()` | L588595 | Reset DNS to 1.1.1.1 on ACTIVE_SERVICE + EN0_SERVICE |
| 21 | `cleanup_network_state()` | L600645 | Nuclear cleanup: disable proxies, flush DNS cache, flush routes, power-cycle Wi-Fi, kill sharing services |
| 22 | `enable_socks_proxy()` | L663682 | Enable system-wide SOCKS5 proxy on macOS |
| 23 | `disable_socks_proxy()` | L684689 | Disable SOCKS5 proxy |
| 24 | `stop_local_dns_proxy()` | L702710 | Kill local Python DNS proxy |
| 25 | `start_local_dns_proxy()` | L713768 | Start Python DNS proxy (UDP:53  TCP:5354) |
| 26 | `force_dns_to_gateway()` | L780820 | Force DNS to a specific IP with 3-retry verification loop |
| 27 | `is_gateway_active()` | L833854 | Check if gateway is running (containers, DNS, SOCKS proxy) |

**Constants:**

| Constant | Value | Line |
|----------|-------|------|
| `LOG_FILE` | `$PWD/output.log` | L8 |
| `PID_FILE` | `/tmp/nullexit-caffeinate.pid` | L105 |
| `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L169 |
| `WARP_WATCHER_PID_FILE` | `/tmp/nullexit-warp-watcher.pid` | L170 |
| `HOST_LEAK_PROBE_PID_FILE` | `/tmp/nullexit-host-leak-probe.pid` | L171 |
| `SOCKS_PROXY_PORT` | `1080` | L661 |
| `PATH` export | `/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH` | L396 |
| Gateway active marker | `/tmp/nullexit-gateway-active.marker` | L89 |
| TUNNEL_FAILED_CLOSED marker | `$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker` | L94 |
| WARP endpoints | `162.159.192.1`, `162.159.193.1` | L535-544 |
```

```
| DNS fallback | `1.1.1.1` | L592 |
| DNS search domain | `ts.net` | L794 |
| `LOCK_FILE` | `/tmp/nullexit-toggle.lock` | L62 |

---

### recover.sh (566 lines, 23KB)

**Functions (7 total):**

| # | Function | Lines | Purpose |
|---|----------|-------|---------|
| 1 | `step()` | L40 | Print bold step header |
| 2 | `ok()` | L41 | Print green success message |
| 3 | `warn()` | L42 | Print yellow warning message |
| 4 | `fail()` | L43 | Print red failure message |
| 5 | `die()` | L44 | Print red error and exit |
| 6 | `run_with_timeout()` | L5674 | Run command with timeout (bash-native, no GNU coreutils) |
| 7 | `restart_tailscaled_daemon()` | L118129 | Restart tailscaled daemon via brew services |

**Constants:**

| Constant | Value | Line |
|----------|-------|------|
| Color codes | `RED`, `GREEN`, `YELLOW`, `BOLD`, `NC` | L37-38 |
| `SCRIPT_DIR` | `$(dirname "${BASH_SOURCE[0]}")` | L50 |
| `PATH` export | `/opt/homebrew/bin:/usr/local/bin:$PATH` | L77 |
| `PID_FILE` | `/tmp/nullexit-caffeinate.pid` | L387 |
| `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L407 |
| `HOST_LEAK_PROBE_PID_FILE` | `/tmp/nullexit-host-leak-probe.pid` | L423 |
| WARP endpoints | `162.159.192.1`, `162.159.193.1` | L329-336 |
| TUNNEL_FAILED_CLOSED marker | `$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker` | L51 |

**Duplicated logic from toggle.sh:**
*(Note: As of July 2026, **ALL** of the below duplicated logic has been successfully extracted to `scripts/common.sh`!)*
- `run_with_timeout()`  simpler subshell version vs toggle's PID-tracking version
- `restart_tailscaled_daemon()`  near-identical (uses warn()/ok() vs echo)
- Active network service detection  inline at L217-233 (toggle has `get_active_service()` function)
- EN0 service resolution  inline at L230-233 (same as toggle L515-516)
- WARP bypass routes  inline at L329-337 (toggle has `add_warp_bypass_routes()`)
- Exit node routing setup  inline at L340-345 (toggle has `setup_exit_node_routing()`)
- DNS cache flushing  L296-302 (same as toggle L611-616)
- Wi-Fi power-cycling  L442-461 (similar to toggle L626-637)
- Proxy disabling  L282-288 (same as toggle L604-608)
- Sharing services reset  L586 (same as toggle L642)
- PID file stop pattern  L387-399, L407-419, L423-435 (same as toggle L157-167, L193-203, L312-322)
- KILL_SWITCH check  L194 (same as toggle L141, L469)
- .gateway_ip reading  L138-142, L209-213 (same pattern as toggle L1080)

---

### setup.sh (410 lines, 18KB)

**Functions (4 total):**

| # | Function | Lines | Purpose |
|---|----------|-------|---------|
| 1 | `step()` | L10 | Print bold blue step header |
| 2 | `ok()` | L11 | Print green success message |
| 3 | `warn()` | L12 | Print yellow warning message |
| 4 | `die()` | L13 | Print red error and exit |

**Constants:**

| Constant | Value | Line |
|----------|-------|------|
| Color codes | `RED`, `GREEN`, `YELLOW`, `BLUE`, `BOLD`, `NC` | L7-8 |
| `RECOMMENDED_TS_VERSION` | `1.98.5` | L71 |
| `ADGUARD_PASSWORD` | `nullexit` | L215 |
| Default .env values | `GATEWAY_RULE_PROFILE=medium`, `GATEWAY_MSS=1120`, etc. | L240-248 |

---

## Part 2: Linux Scripts

### toggle.sh  Linux path (unified into the root cross-platform script)

The Linux toggle is no longer a separate file. `scripts/toggle-linux.sh` was
merged into the repo-root `toggle.sh`, which dispatches on `$OSTYPE`: an early
`if [[ "$OSTYPE" != darwin* ]]` block runs the self-contained Linux toggle
```

```
(shuf / ss / nmcli / ip / systemd) and exits before the macOS body. The shared
DNS/proxy/daemon primitives it used were already moved to `common.sh` in the
`313dc32` unification, so both branches call them by the same names.

---

### recover.sh  Linux path (unified into the root cross-platform script)

Linux recovery is no longer a separate file. `scripts/recover-linux.sh` was
merged into the repo-root `recover.sh`, which dispatches on `$OSTYPE`: the Linux
branch runs a self-contained teardown (resolvectl / nmcli / ip) and exits before
the macOS dual-mode body. All formatting/timeout helpers come from `common.sh`.

---

### scripts/setup-linux.sh (416 lines, 18KB)

**Functions (4 total):**

| # | Function | Line | Purpose |
|---|----------|------|---------|
| 1 | `step()` | L10 | Print formatted step header |
| 2 | `ok()` | L11 | Print green success message |
| 3 | `warn()` | L12 | Print yellow warning message |
| 4 | `die()` | L13 | Print red error + exit 1 |

**~95% identical to macOS setup.sh.** Only differences:
1. Desktop shortcuts (L366-391 Linux `.desktop` files vs macOS `.app` via `osacompile`)
2. Final instructions text
3. .env template: macOS adds `GATEWAY_USE_EXIT_NODE`, `WARP_FAIL_THRESHOLD`, `HOST_LEAK_PROBE`, `KILL_SWITCH`
     Linux omits these

---

## Part 3: Utility Scripts

### scripts/diagnose-host-leak.sh (722 lines, 38KB)

**Purpose:** Comprehensive host-routing diagnostic. 8 checks classify host IP leaks. Supports `--fix` and `--
     watch` modes.

**Duplicated logic:**
- `resolve_active_service()`  identical pattern as toggle.sh, recover.sh, fix-docker-bridge-collision.sh
- `run_with_timeout()`  reimplemented timeout
- Color codes  RED/GREEN/YELLOW/BOLD/NC with tty detection
- WARP check via `cloudflare.com/cdn-cgi/trace`
- `export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"`
- .gateway_ip reading pattern

### scripts/fix-docker-bridge-collision.sh (402 lines, 18KB)

**Purpose:** Fixes Docker bridge subnet collisions with `10.200.1.0/24`.

**Duplicated logic:**
- `resolve_active_iface()`  same interface detection as diagnose-host-leak.sh
- Color codes, step()/ok()/warn()/fail()/die() formatting helpers

### scripts/fix-ssh-delay.sh (63 lines, 2.7KB)

**Purpose:** One-shot fix for ~20s SSH delay. Standalone, minimal duplication (uses hardcoded / instead of
     color functions).

### scripts/host-leak-probe.sh (110 lines, 5.5KB)

**Purpose:** Sub-second host-egress leak prober at 300ms intervals.

**Duplicated logic:**
- WARP check via cdn-cgi/trace (same curl + awk as toggle.sh WARP Watcher)
- Color codes (hardcoded; status updates and warnings are conditioned on TTY detection to prevent log pollution
     )
- LOG_FILE = `$PROJECT_ROOT/output.log`

### scripts/reinstall-host-tailscale.sh (740 lines, 31KB)

**Purpose:** Complete uninstall + reinstall of host Tailscale.

**Duplicated logic:**
- `with_timeout()`  yet another bash timeout implementation (polling-based, different from diagnose and toggle
     versions)
- Color codes + logging functions
- System Extension detection logic (similar to diagnose-host-leak.sh)
```

```
234
235 ### scripts/routing-fix.sh (266 lines, 10.5KB)
236
237 **Purpose:** Runs INSIDE warp container. Maintains routing table 200, FORWARD rules, country blocking, IP
      blocklists, tunnel health checks.
238
239 **Duplicated logic:**
240 - (Fixed) WARP_ENDPOINT defaults are now centralized in common.sh
241 - WARP health check via cdn-cgi/trace (another instance)
242 - iptables dual-backend pattern (for-loop over `iptables iptables-legacy`)
243
244 ### scripts/watcher.sh (269 lines, 13KB)
245
246 **Purpose:** Long-running launchd daemon for post-wake/post-roam recovery.
247
248 **Key function  `detect_lan_p2p_mode()`:**
249 Added July 10, 2026. Called on startup and on every network state change (Listener 2 NET: events). Determines
      whether the current Wi-Fi network safely allows direct Tailscale LAN P2P connections:
250 1. **WPA2-Enterprise check:** `system_profiler SPAirPortDataType`  reads the `Security:` field for the current
      association. Enterprise = 802.1x = always AP-isolated  sets `allow=false`.
251 2. **AP isolation probe:** `route -n get 1.1.1.1`  finds default gateway IP, then `ping -c 3 -t 1 -q "$gw"`  if
       0 responses on a non-empty network, AP isolation is active  sets `allow=false`.
252 3. **Output:** Writes `'true'` or `'false'` to `.lan_p2p_detected` in the repo root.
253 4. **Consumer:** `routing-fix.sh` reads this file every 30 seconds and enforces/relaxes the RFC1918 DROP rule
      accordingly  no restart required.
254
255 **Duplicated logic:**
256 - PATH export (similar to toggle.sh/recover.sh)
257 - Marker file pattern (`/tmp/nullexit-gateway-active.marker`)
258
259 ### scripts/dns-proxy.py (92 lines, 2.9KB)   Clean, standalone
260
261 ### scripts/logger.py (144 lines, 5.9KB)   Clean, standalone
262
263 ---
264
265 ## Part 4: Cross-Cutting Issues
266
267 ### Functions Identical Across macOS  Linux
268
269 | Function | macOS toggle | macOS recover | macOS setup | Linux toggle | Linux recover | Linux setup |
270 |----------|:-----------:|:------------:|:-----------:|:------------:|:-------------:|:----------:|
271 | `run_with_timeout()` |  |  |  identical |  identical |  |
272 | `stop_local_dns_proxy()` |  |  |  identical |  |  |
273 | `start_local_dns_proxy()` |  |  | ~95% similar |  |  |
274 | `is_gateway_active()` |  |  | ~80% similar |  |  |
275 | `step/ok/warn/die` |  |  |  identical |  identical |
276
277 ### Platform-Adapted Functions (same logic, different OS commands)
278
279 | Function | Linux Command | macOS Command |
280 |----------|--------------|---------------|
281 | `get_active_service()` | `ip route get 1.1.1.1 \| awk '{print $5}'` | `route get default` + `networksetup -
      listnetworkserviceorder` |
282 | `reset_dns()` | `resolvectl revert` | `networksetup -setdnsservers` |
283 | `cleanup_network_state()` | `ip link set dev down/up` | `networksetup -setairportpower off/on` + proxy
      disable |
284 | `force_dns_to_gateway()` | `resolvectl dns` + `resolvectl domain ~ts.net` | `networksetup -setdnsservers` +
      `-setsearchdomains ts.net` |
285 | `start_sleep_prevention()` | `systemd-inhibit --what=sleep` | `caffeinate -i` |
286 | `enable/disable_socks_proxy()` | Stub (no-op on Linux) | `networksetup -setsocksfirewallproxy` |
287 | DNS cache flush | `resolvectl flush-caches` | `dscacheutil -flushcache` + `killall -HUP mDNSResponder` |
288
289 ### Features Missing From Linux Scripts
290
291 These exist in `toggle.sh`'s macOS branch but NOT its Linux branch:
292
293 1. `write_gateway_active_marker()` / `clear_gateway_active_marker()`
294 2. `restart_tailscaled_daemon()`
295 3. `start_warp_watcher()` / `stop_warp_watcher()`  WARP liveness monitor
296 5. `add_warp_bypass_routes()` / `remove_warp_bypass_routes()` / `setup_exit_node_routing()`
297 6. `LOG_FILE` structured logging (timestamps to output.log)
298 7. MSS clamping injection (step 4c)
299 8. Rule count reporting
300 9. `--post-wake` mode in recover
301 10. Container teardown on failure in cleanup_handler
302
303 ---
304
305
306
```

```
## Part 5: Constant Duplication Map

| Value | Files |
|-------|-------|
| `162.159.192.1` (WARP endpoint) | docker-compose.yml, routing-fix.sh (now dynamically loaded via .env/common.sh!) |
| `5335` (AdGuard DNS port) | docker-compose.yml, AdGuardHome.yaml, post-rules.txt |
| `5354` (mapped DNS port) | docker-compose.yml, dns-proxy.py |
| `41642` (Tailscale WireGuard) | docker-compose.yml, routing-fix.sh, post-rules.txt |
| `1080` (SOCKS5 port) | docker-compose.yml, toggle.sh |
| `1120` (MSS clamp) | post-rules.txt, .env |
| `admin:nullexit` (AdGuard creds) | AdGuardHome.yaml (bcrypt) |
| `10.200.1.0/24` (Docker subnet) | (Fixed) no longer duplicated |
| `America/New_York` (timezone) | docker-compose.yml (4 services) |
| `/tmp/nullexit-caffeinate.pid` | (Fixed) centralized in common.sh |
| `/tmp/nullexit-dns-watcher.pid` | (Fixed) centralized in common.sh |
| `/tmp/nullexit-host-leak-probe.pid` | toggle.sh, recover.sh |
| `/opt/homebrew/bin:...` (PATH) | (Fixed) centralized in common.sh |

---

## Part 6: Refactoring Outcome (Implemented July 2026)

**Decision: "Scripts-only" architecture**
Instead of introducing a new `lib/` directory, all shared code was moved into the existing `scripts/` directory to keep the project hierarchy flat and simple.

### 1. `scripts/common.sh` (288 lines, 10KB)
Extracts the fundamental primitives and networking logic used across orchestrator scripts:
- **Formatters:** `step()`, `ok()`, `warn()`, `fail()`, `die()`
- **Timeout Wrapper:** `run_with_timeout()` (Canonical version with PID tracking, graceful SIGTERM, and SIGKILL fallback)
- **Daemon Lifecycle:** `restart_tailscaled_daemon()`, `stop_pidfile_daemon()` (collapsed caffeinate, DNS watcher, WARP watcher, and host leak prober teardown logic)
- **Network Interface/Routing:** `get_active_service()`, `get_en0_service()`, `bounce_wifi_interfaces()`, `add_warp_bypass_routes()`, `remove_warp_bypass_routes()`
- **Proxy/DNS Cleanup:** `disable_all_proxies()`, `flush_dns_cache()`, `reset_sharing_services()`
- **Configuration Helpers:** `read_adguard_ip()`, `is_kill_switch_enabled()`, `get_warp_endpoint_1()`, `get_warp_endpoint_2()` (which centralizes the hardcoded 162.159.192.1 / .193.1 into .env logic)

### 2. `scripts/setup-common.sh` (~350 lines)
Extracts the heavy, duplicated installation logic shared between `setup.sh` (macOS) and `scripts/setup-linux.sh`:
- Docker and Colima prerequisite checks
- `wgcf` binary download and WARP key generation
- `.env` configuration file generation
- Interactive Tailscale Auth Key and AdGuard Rule Profile prompts

### Sourcing Pattern
- macOS root scripts (`toggle.sh`, `recover.sh`, `setup.sh`) source via:
  `source "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/common.sh"`
- Scripts under `scripts/` (`watcher.sh`, `diagnose-host-leak.sh`, etc.) source via:
  `source "$(dirname "${BASH_SOURCE[0]}")/common.sh"`

## Agent Verbs / Protocols

### `/sweep` (or `sweep the gateway`)
When the user invokes this verb, the agent MUST perform a full end-to-end connectivity, health, security, and performance audit of the gateway:
1. **WARP Validation:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp=` (must equal `on`).
2. **Tor Proxy & Control Validation:**
   - Source the dynamically generated ports from `/tmp/nullexit-ports.env` (or identify them using `docker compose ps`).
   - **SOCKS5 Check:** Run `curl -s --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip` to confirm the SOCKS5 proxy successfully connects and routes through the Tor network.
   - **Control Port Check:** Query the Tor Control Port using netcat: `echo "PROTOCOLINFO" | nc 127.0.0.1:$TOR_CONTROL_PORT`. Verify it returns a standard `250-PROTOCOLINFO` response, confirming the Control Port is active and communicating.
   - **Circuit & Node Lookup:** Query the Tor Control Port for the active path (`GETINFO circuit-status`). For each active built circuit, parse the relay fingerprints, lookup their IP addresses (`GETINFO ns/id/<fingerprint>`), and resolve their country codes using Tor's local GeoIP database (`GETINFO ip-to-country/<IP>`). Include the list of active circuits, showing the role (Guard/Middle/Exit), nickname, IP, and country of each hop in the final sweep report.
   - **Transparent Onion Validation:** Run `nslookup duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion` to verify it resolves to an IP within `198.18.0.0/15` benchmark range. Verify transparent dark web routing by running `curl -sI -H "Host: duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion" http://<resolved-ip>/` (should return HTTP 301/200).
3. **DNS Leak Validation:**
   - Audit the upstream resolver currently utilized by the host. Run `curl -s https://edns.ip-api.com/json` to fetch details of the DNS resolver processing the client's requests.
```

```
   - Verify that the resolver IP and organization returned belong to Cloudflare (e.g., AS13335) or match the WARP tunnel's exit range, and that **no** DNS traffic leaks to the user's raw physical ISP DNS.
4. **Performance Benchmark Audit:**
   - Run the python benchmark script `python3 benchmarks/test_load.py`.
   - Report the overall load speed and the ratio of successfully fetched vs. blocked/failed subresources for the targeted domains to ensure ad-blocking and MTU encapsulation are performing correctly.
5. **Tailscale Mesh Validation:**
   - Run `ping -c 1 100.100.21.8` to ensure the gateway is reachable.
   - Run `tailscale status` to identify all active/online nodes in the mesh.
   - For every online peer returned in the status output, execute a ping check (`ping -c 1 <peer-ip>`) to verify bidirectional connectivity and confirm healthy routing to all active mesh devices.
6. **Log Audit:** Run `tail -n 100 output.log | grep -iE "(error|fail|warn)"` to catch any underlying component failures or warnings.

### `/latex` (or `generate latex`)
When the user invokes this verb, the agent MUST run `python3 -I scripts/generate_tex.py` to regenerate the documentation source code (`nullexit_unified.tex`). The agent MUST NOT attempt to compile the `.tex` file into a PDF (the user handles PDF compilation locally).
```

4 devref.md

```
# nullexit  Development Reference & Resolved Issues

> **Last updated:** July 12, 2026
> **Purpose:** Provide any LLM or developer with complete project understanding, debugging history, and resolved issues so they can make informed changes without re-reading every file.
> **Diagrams:** See [`diagrams.md`](./diagrams.md) for system architecture, toggle.sh flowcharts, monitoring layer, traffic sequence, recover.sh decision tree, and the full failureself-healing map.

This reference is organized into four parts:
> - **Part I  Architecture & Reference** (18)
> - **Part II  Design Decisions & Threat Model** (914)
> - **Part III  Incident Log & Resolved Issues** (15)
> - **Part IV  Observations, Changelog & TODO** (1618)

---

# Part I  Architecture & Reference

## 1. What Is nullexit?

nullexit is a **Tunnel-in-Tunnel VPN gateway** that chains **Tailscale** (mesh VPN) through **Cloudflare WARP** (exit VPN). It runs on a macOS host inside Docker containers managed by **Colima** (lightweight Linux VM). The goal: any device on the user's Tailscale mesh can use this gateway as an exit node, and all traffic exits through Cloudflare WARP  achieving double encryption, ISP invisibility, and network-wide ad-blocking via **AdGuard Home**.

### Traffic Flow
```
Phone/Laptop Tailscale mesh (WireGuard) Mac host Docker container
   Tailscale decrypts Gluetun re-encrypts WARP tun0 Cloudflare Internet
```

### SOCKS5 Failover Path (if exit node is unavailable)
```
DNS: Browser host 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP DNS
TCP: Apps macOS SOCKS5 proxy (127.0.0.1:1080) container tun0 WARP Internet
```

---

## 2. Architecture

### Containers (all share `warp`'s network namespace via `network_mode: service:warp`)
| Container | Image | Role |
|-----------|-------|------|
| `warp` | `qmcgaw/gluetun:v3.41.1` | Gluetun WireGuard client  Cloudflare WARP. Owns the network namespace. Strict firewall. |
| `tailscale` | `tailscale/tailscale:v1.98.4` | Advertises as exit node on the mesh. Also provides the built-in SOCKS5 proxy fallback via `TS_SOCKS5_SERVER=:1080`. |
| `routing-fix` | `alpine:3.20` | Sidecar that maintains routing tables + iptables rules every 30 seconds. |
| `rule-compiler`| `golang:1.22-alpine` / `alpine:3.20` | One-shot startup container that compiles 500k+ DNS block rules in <2s and immediately exits. |
| `adguardhome` | `adguard/adguardhome:v0.107.77` | DNS sinkhole for ads/trackers. Listens on port 5335. Upstream DNS: `127.0.0.1:53` (through WARP). |

```

```
### Port Mappings (on host via Colima)
| Host Port | Container Port | Protocol | Purpose |
|-----------|---------------|----------|---------|
| 5354 | 5335 | TCP+UDP | AdGuard DNS |
| (Tailscale IP):3000 | 3000 | TCP | AdGuard web UI (Internal to Mesh) |
| 41641 | 41641 | UDP | Tailscale WireGuard direct |
| 1080 | 1080 | TCP | SOCKS5 proxy |

### Host Components
- **Colima VM**  Linux VM running Docker (600MB RAM + 400MB swap)
- **tailscaled**  Standalone Tailscale daemon via `brew install tailscale` + `brew services start tailscale` (NOT the GUI `.app`)
- **macOS network settings**  DNS, SOCKS proxy, routing all manipulated by `toggle.sh`

---

## 3. Key Files

| File | Lines | Purpose |
|------|-------|---------|
| `toggle.sh` | ~2109 | **Main script (macOS + Linux).** Dispatches on `$OSTYPE`  an early Linux block (shuf/ss/nmcli/ip, systemd) runs and exits before the macOS body (jot/netstat/networksetup, launchd/pf). Detects state, toggles gateway ON/OFF. Handles DNS hijacking, Tailscale exit node, SOCKS proxy, sleep prevention, timeouts, cleanup. |
| `setup.sh` | 406 | **One-time setup.** Installs deps (Docker, Tailscale, wgcf), generates WARP keys, writes `.env`, configures AdGuard via API, starts containers. |
| `docker-compose.yml` | 194 | Service definitions for all 5 containers. |
| `scripts/routing-fix.sh` | 201 | Maintains routing tables (table 200 for SOCKS5, table 52 for CGNAT) and FORWARD iptables rules in both nftables and legacy backends. Loads and enforces the IP blocklist via `ipset`. Runs in a 30-second loop. |
| `post-rules.txt` | 18 | Gluetun iptables rules loaded at container start. FORWARD accept for tailscale0, NAT MASQUERADE on tun0, DNS redirect to port 5335, IPv6 drop, TCP MSS clamping. |

| `scripts/dns-proxy.py` | 91 | Local DNS proxy. UDP:53  TCP:5354 with proper 2-byte length prefix. Fallback when exit node is unavailable. |
| `scripts/rule-compiler/` | ~800 | Unified threat compiler (ported to Go from Python). **DNS pipeline:** fetches remote blocklists, deduplicates against AdGuard native filters, optimizes subdomains (~50% reduction), compiles to AdGuard syntax. **IP pipeline:** fetches threat-intel feeds (Spamhaus, Feodo, ET, CINS), strips private/Tailscale ranges, outputs atomic `ipset` restore file. Built dynamically in Docker multi-stage build. |
| **`adguard/work/`** |  | **Bind-mounted container data folder. Off-limits from outside Docker  READ-ONLY-EXCEPT-VIA-`sync-rules`. See README 6.** |
| `recover.sh` | ~742 | Nuclear recovery script (**macOS + Linux**  dispatches on `$OSTYPE`; the Linux branch runs a self-contained teardown via `resolvectl`/`nmcli`/`ip` and exits, macOS falls through to the dual-mode body). Gain: use `--post-wake` (macOS) for non-destructive refresh after sleep/wake or Wi-Fi roam keeps the gateway live while still re-hijacking DNS + refreshing Tailscale exit-node + force-recreating warp if unhealthy. Resets DNS on ALL services, disables proxies, flushes routes, stops containers, kills sleep prevention, power-cycles Wi-Fi. Cryptographically verified. |
| `scripts/diagnose-host-leak.sh` | 580+ | **One-shot host-routing diagnostic.** Runs 8 checks (Tailscale, SOCKS5, CDN-cgi/trace, IPv6 leak, default route, output.log hints, gateway IP, Tailscale System Extension conflict), classifies into ONE of three known leak scenarios (A=SOCKS5 fallback, B=IPv6 leak, C=route freeze) or OK, writes a timestamped report to `host-leak-diagnostic-<UTC>.txt`, and prints a ready-to-run fix on stdout. Pass `--fix` to apply the matched remediation and re-verify egress. Pass `--watch` (or `--watch 30`) to run a full baseline then continuously monitor warp/IPv6/default-route every N seconds, alerting on any state change. See 15.6.1. |
| `scripts/host-leak-probe.sh` | ~110 | **Continuous sub-second host-egress prober.** Launched automatically by `toggle.sh` when `HOST_LEAK_PROBE=true` in `.env` (default). Polls `cdn-cgi/trace` every 300ms directly from the host NIC via `curl`  not via `docker compose exec`  so it detects flash-leaks invisible to the in-container WARP Watcher. Logs only on state change (`LEAK`, `ROTATE`, `HOST-PROBE failed/timeout`) to `output.log`. curl errors (`2>>output.log`) also land there. Stopped gracefully (SIGTERM) by `toggle.sh` STOP path, `cleanup_handler`, and `recover.sh`. See 15.6.1 (Host Leak Probe deep dive). |
| `scripts/fix-docker-bridge-collision.sh` | ~290 | **One-shot fix for 15.2.3 (Docker bridge subnet collision).** Detects Docker's `172.17.0.0/16` bridge colliding with the host Wi-Fi subnet (the symptom that produces `default 172.17.0.1 UGScg en0` in `netstat -rn`), picks a non-colliding `docker.bip` from a small candidate set, idempotently edits `~/.colima/default/colima.yaml` with a timestamped backup, restarts Colima, rebinds Wi-Fi DHCP, verifies the new default route is no longer Docker's bridge, and re-runs `toggle.sh` + `scripts/diagnose-host-leak.sh`. Supports `--dry` and `--skip-toggle`. |
| `scripts/watcher.sh` | 194 | **Long-running post-wake + post-roam daemon.** Launched by launchd LaunchAgent; subscribes to `com.apple.powermanagement` events via `log stream` and to network-state changes via `scutil n.watch`. Both fire `recover.sh --post-wake`. Single-instance lock + SCRIPT_DIR-relative resolve of `RECOVER`. See 15.5.1. |
| `scripts/common.sh` | ~40 | **Shared script primitives.** Centralizes formatted echo statements (`step`, `ok`, `warn`, `die`) and the safe background `run_with_timeout` execution wrapper. Sourced by all orchestrator scripts across macOS and Linux. |
| `scripts/setup-common.sh` | ~350 | **Shared installation logic.** Centralizes logic for verifying Docker, installing Tailscale/wgcf, generating `.env`, and prompting for auth keys. Sourced by `setup.sh` and `setup-linux.sh`. |
| `scripts/setup-linux.sh` | ~80 | **Linux sibling of `setup.sh`.** Sources `setup-common.sh`. Distro detection (apt/dnf/pacman), installs Linux-specific dependencies, creates `.desktop` shortcuts. |
| `scripts/Toggle-Gateway.bat` | ~300 | **Windows Toggle helper.** Moved from root to `scripts/`. Helps invoke the framework on Windows if applicable. |
```

```
| `scripts/reinstall-host-tailscale.sh` | ~740 | **Nuke-and-reinstall tool for host Tailscale.** Completely purges Tailscale, its settings, macOS Background Items (`.btm` database), and stale System Extensions. Wipes ALL Login Items via `sfltool reset-login-items` (requires sudo). Useful for crisis recovery but destructive to other login items. |
| `scripts/fix-ssh-delay.sh` | ~100 | **Fixes SSH connection delays.** One-shot script to address SSH delays caused by DNS or routing issues, useful in crisis situations. |
| `scripts/logger/` | ~100 | **Internal logger piped from `scripts/routing-fix.sh`** inside the `warp` container. Ported to Go from Python. Built via Docker multi-stage build. |
| `black_list.txt` | 71 | Custom domains to block (ads, trackers, telemetry). Supports `$important` modifier. |
| `white_list.txt` | 222 | Domains to force-allow (YouTube, Apple services, etc.). Always wins over blocks. |
| `.env` | ~14 | WARP WireGuard keys, Tailscale auth key, rule profile. **Contains secrets & NULLEXIT_SEED.** |
| `.gateway_ip` | 1 | Static gateway Tailscale IP (fallback for dynamic resolution). |
| `scripts/unlock-files.sh` | ~20 | **One-shot stale-permission fix.** Uses atomic rename (`cp` + `mv`) to replace locked inodes (mode `000`/`0444` from old `chmod` decisions) with fresh writable ones  no `chmod` called. |
| `scripts/crypto.sh` | ~50 | **Cryptographic integrity enforcement.** Uses `NULLEXIT_SEED` from `.env` to sign core bash scripts (HMAC-SHA256) and strictly verify them on start to prevent manipulation. Pass `--sign` or `--verify`. |
| `scripts/pf.conf` | 35 | **macOS native firewall ruleset.** Enforces the default-deny kill-switch, allows local LAN / Tailscale traffic, and applies TCP MSS Clamping (`max-mss 1160`) to prevent WireGuard MTU fragmentation. |

---

## 4. toggle.sh Flow

### State Detection
`is_gateway_active()` returns true if EITHER containers are running OR host DNS is not `1.1.1.1`. This dual check prevents the script from starting when it should stop (e.g., containers crashed but DNS is still hijacked).

### START Path (gateway OFF  ON)
1. **Reset DNS to 1.1.1.1**  Prevents deadlocks during startup
2. **Disconnect host Tailscale**  `tailscale down` (prevents exit-node routing during container boot)
3. **Compile DNS rules**  `docker compose run --rm rule-compiler` (Go binary inside multi-stage Docker build)
4. **Boot Colima VM** (if not running)  via `colima_start_until_ready()`, which returns as soon as Docker is reachable (~12s) instead of blocking on Colima's ~2-min post-provision hang (15.8.7). The networking mode is **gated on the LAN-P2P signal** (`.lan_p2p_detected`, falling back to `TAILSCALE_ALLOW_LAN_P2P` in `.env`): a trusted home network requests `--network-address --network-mode bridged` (or `shared` on WPA2-Enterprise), while the common AP-isolated / untrusted case boots plain fast user-mode networking (`colima start --memory 0.6 --vm-type vz`).
   - **Why Bridged Networking (when enabled)?** The `--network-mode bridged` and `--network-address` flags force the VM to receive its own IP directly from your local router's DHCP server, placing it on the same physical subnet as the host. This bypasses macOS network isolation and is what lets **external LAN devices** (phones, laptops on the same LAN), mDNS, and local service discovery reach across the container boundary. It requires the `socket_vmnet` helper and is only worthwhile on a trusted LAN, which is why it is gated rather than always requested. **It is *not* required for direct hostcontainer P2P**  that path (the gateway its own Mac) works even in user-mode networking via the routing-fix `.host_ips` RETURN-rule carve-out; see 15.3.5 and 15.8.7.
   - **Enterprise Wi-Fi Fallback & Dynamic Roaming:** `toggle.sh` dynamically queries `system_profiler SPAirPortDataType` on boot. If it detects a **WPA2-Enterprise** connection (802.1X), it forces Colima to fall back to `--network-mode shared`. This prevents aggressive network intrusion systems on corporate/university Wi-Fi switches from terminating the host's Wi-Fi connection due to "MAC spoofing" (detecting multiple MAC addresses originating from a single authenticated port). P2P connections are impossible on these networks anyway due to AP Client Isolation.
   - **Dynamic Roam Downgrading:** `toggle.sh` writes its selected network mode to `/tmp/nullexit_colima_mode.txt`. When roaming between networks (e.g., Home to University), `recover.sh` actively checks this state against the new network's security requirement (`.lan_p2p_detected`). If a mismatch is detected (e.g., bridged Colima on an Enterprise network), `recover.sh` automatically forces a full `toggle.sh --restart` in the background to safely transition the Virtual Machine's network mode and protect the host from de-authentication.
5. **Configure VM swap**  400MB swap file inside the VM to prevent OOM
6. **Clean corrupted AdGuard config**  Remove empty `AdGuardHome.yaml`
7. **Start containers**  `docker compose up -d`
8. **Wait for gateway Tailscale**  Poll `tailscale status` for "offers exit node" (up to 60s, abort at 40 consecutive NoState)
9. **Resolve gateway IP**  From `.gateway_ip` or `docker compose exec tailscale tailscale ip -4`
10. **Connect host to mesh**  Verify `tailscaled` is running (auto-start if needed), then `tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true --exit-node=` (`--accept-routes=true` must be explicit `--reset` reverts it to `false` otherwise, silently preventing the default route from switching to `utun*`)
11. **Pre-flight checks**  [1/3] `tailscale ping` gateway, [2/3] `dig +tcp` AdGuard DNS, [3/3] WARP container internet
12. **If all pass**  Set exit node + hijack DNS to gateway IP (single server, no fallback)
13. **If any fail**  Enable SOCKS5 proxy + local DNS proxy as fallback
14. **Final DNS re-force**  `force_dns_to_gateway` called again after exit-node transition (tailscaled clobbers DNS briefly)
15. **Enable sleep prevention**  `caffeinate -i` in background to prevent idle sleep

### STOP Path (gateway ON  OFF)
1. Disconnect host Tailscale (`tailscale down`)
```

```
2. `docker compose down -t 5`
3. Reset DNS to 1.1.1.1
4. Stop local DNS proxy
5. **Stop sleep prevention**  Kill caffeinate process
6. Full network state cleanup (proxies off, DNS cache flush, route flush, Wi-Fi power cycle)

### Safety Mechanisms
- `run_with_timeout()`  Pure-bash watchdog for every system command (15s120s)
- `cleanup_handler()`  Trap on ERR/INT/TERM/HUP restores DNS to 1.1.1.1, kills caffeinate, and tears down Tailscale
- `force_dns_to_gateway()`  Sets DNS and VERIFIES via `networksetup -getdnsservers` (3 attempts with backoff)
- `SUCCESS_RUN` flag  Cleanup only runs if script didn't complete successfully
- `start_sleep_prevention()` / `stop_sleep_prevention()`  Manages `caffeinate -i` via PID file at `/tmp/nullexit-caffeinate.pid`

---

## 5. Routing & Firewall Architecture (Inside Container)

### IP Rules (`ip rule show`)
| Priority | Match | Table | Purpose |
|----------|-------|-------|---------|
| 99 | `to 100.64.0.0/10` | 52 | CGNAT range  Tailscale routes (injected by routing-fix) |
| 100 | `from 172.18.0.2` | 200 | Container's own traffic (SOCKS5 proxy)  tun0 |
| 101 | all | 51820 | Gluetun's WireGuard fwmark  tun0 |

### Table 200 (SOCKS5 proxy traffic)
- `162.159.192.1 via $DOCKER_GW dev eth0`  WARP endpoint exception (prevents tunnel loop)
- `default dev tun0`  Everything else through WARP

### iptables (TWO separate stacks!)
- **nftables backend** (`iptables`)  Used by Gluetun's `post-rules.txt`. FORWARD policy DROP.
- **legacy backend** (`iptables-legacy`)  Used by Tailscale's `ts-forward`. FORWARD policy ACCEPT.
- Both stacks apply to every packet. FORWARD RELATED,ESTABLISHED must exist in BOTH.

### post-rules.txt (loaded by Gluetun)
```
FORWARD: ACCEPT for tailscale0 in/out
INPUT: ACCEPT for tailscale0
NAT POSTROUTING: MASQUERADE on tun0
MANGLE: TCP MSS clamp to 1180
NAT PREROUTING: Redirect DNS (53) 5335 on tailscale0 and eth0
ip6tables: DROP FORWARD on tailscale0
```

---

### 5.1 Tailscale Local Network Discovery (DERP Bypass)
By default, Gluetun's strict NAT swallows all of Tailscale's UDP hole-punching packets, forcing Tailscale to fall back to incredibly slow DERP relay servers (often adding 500ms+ latency).

To fix this, we use `FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8` in `docker-compose.yml` (on the `warp` container). This creates `ip rule` bypasses (Priority 99, Table 199) that route all packets destined for private IP ranges *outside* the WARP tunnel directly to the host's `eth0`.
- **172.16.0.0/12** is especially critical because many modern Wi-Fi routers (and Docker itself) assign IPs in this block (e.g., `172.17.x.x`).
- This bypass allows Tailscale's UDP packets to reach devices on the same Wi-Fi directly, establishing a fast peer-to-peer connection while the actual web traffic inside the tunnel remains fully routed through WARP.

---

### 5.2 Firewalling & Per-Device Access Control
Because every mesh device's traffic passes through the `warp` container's network namespace, this namespace serves as the ultimate choke point for network-wide firewall rules.

**Where to Put Rules**
The right place to add firewall rules is `scripts/routing-fix.sh`. It runs a 30-second loop inside the namespace and already handles re-injecting rules that Gluetun resets on reconnect. **Do not use `post-rules.txt` for dynamic rules**, as Gluetun flushes and resets it upon VPN reconnects.

**The Dual iptables Backend Constraint (CRITICAL)**
The container runs two simultaneous iptables backends: `iptables` (for the nftables backend used by Gluetun) and `iptables-legacy` (for Tailscale). Every rule **must** be added to both stacks, or it will silently fail for certain traffic.

**Avoiding Duplication**
Because `scripts/routing-fix.sh` runs in a loop, you must always check for a rule's existence with `-C` before appending with `-A`. Otherwise, your rules will duplicate every 30 seconds.

**Per-Device Identification & Filtering**
```

```
Each mesh device has a stable `100.x.x.x` Tailscale IP, which is visible as the `src` IP in the `FORWARD` chain. This is how you identify devices for filtering.
* **Domain-level:** AdGuard Home (port 3000, credentials `admin/nullexit`) provides a REST API that supports per-client blocking rules based on their Tailscale IP.
* **IP-level:** Use `iptables` in `scripts/routing-fix.sh` (e.g., `iptables -I FORWARD -s 100.x.x.x -d <blocked_ip> -j DROP`).
* **Country-level:** Add `ipset` to the `routing-fix` apk install step in `docker-compose.yml`, and load CIDR ranges from `ipdeny.com` into an ipset, then block that set in the `FORWARD` chain. See 5.3 (Geo-IP Blocking) for the full configuration flow.

### 5.3 Geo-IP Blocking
To block traffic to and from specific countries, the system dynamically downloads country-specific IP ranges from [ipdeny.com](http://www.ipdeny.com/ipblocks/data/countries/) and drops them via iptables `FORWARD` rules.

To add a new country to the blocklist:
1. Find the 2-letter ISO country code (e.g., `cn` for China, `ru` for Russia, `il` for Israel, `kp` for North Korea).
2. Open your `.env` file.
3. Update the `BLOCKED_COUNTRIES` variable: e.g., `BLOCKED_COUNTRIES="kp il cn ru"`.
4. Run `toggle.sh` to restart the gateway and apply the new environment variables to the routing-fix container.

> [!NOTE]
> **Limitations of Geo-IP Blocking:** IP-based blocking is highly effective at blocking servers, direct apps, and raw infrastructure physically located and registered in a specific country (e.g., `www.gov.il`). However, it will **not** block high-profile websites (e.g., `jpost.com`) that route their traffic through global Content Delivery Networks (CDNs) like Google Cloud, Cloudflare, or Fastly. Because the CDN's IP addresses are registered in the US or globally, the network traffic physically never routes to the blocked country, allowing it to bypass the Geo-IP firewall.

## 6. macOS Quirks to Remember

1. **Colima SSH tunnel is TCP-only**  UDP ports (like 5354/udp) are NOT accessible from host. Use `dig +tcp` or the Python DNS proxy (UDPTCP converter).
2. **`docker compose exec -T` injects `\r`**  Always `tr -d '\r'` when capturing output for comparisons or `networksetup` commands.
3. **`brew services` can get stuck**  Fix with `sudo brew services restart tailscale`. Common state: `brew services list` shows `started` but `tailscale status` shows "Tailscale is stopped."
4. **No `timeout` command**  macOS lacks GNU `timeout`. Use the `run_with_timeout()` bash function.
5. **`sudo -v` prompts even with NOPASSWD**  Use `sudo -n` (non-interactive) everywhere.
6. **DNS is scoped to en0**  macOS scutil DNS resolver is scoped to en0 (Wi-Fi). DNS changes on other interfaces (like USB Ethernet) are ignored unless en0's service is also updated. The script always sets DNS on BOTH `$ACTIVE_SERVICE` and `$EN0_SERVICE`.
7. **tailscaled clobbers DNS during exit-node transition**  `force_dns_to_gateway` is called a second time at the end of the ENABLE branch to counteract this.
8. **`tailscale up --reset`** resets all unspecified flags to defaults  Every `--reset` call MUST also pass `--accept-dns=false` explicitly or tailscaled re-enables DNS management.
9. **`docker compose ps` CWD pitfall**  `docker compose` commands require a `docker-compose.yml` in the current working directory. Use `docker ps` for raw container listing from any CWD; `docker compose ps` requires the project directory.

---

## 7. Environment Variables

### `.env` File
| Variable | Purpose |
|----------|---------|
| `WIREGUARD_PRIVATE_KEY` | WARP WireGuard private key (from `wgcf generate`) |
| `WIREGUARD_PUBLIC_KEY` | WARP WireGuard public key |
| `WIREGUARD_ADDRESSES` | WARP WireGuard address (e.g., `172.16.0.2/32`) |
| `TS_AUTHKEY` | Tailscale auth key for the container |
| `NULLEXIT_SEED` | 256-bit seed for HMAC-SHA256 script integrity signing (generated by `setup.sh`) |
| `GATEWAY_RULE_PROFILE` | Rule compilation tier: `light`, `medium`, `heavy` |
| `GATEWAY_BYPASS_PING` | (Optional) `true` to proceed even if pre-flight checks fail |
| `GATEWAY_USE_EXIT_NODE` | (Optional) `false` to skip exit node, DNS-only mode |
| `GATEWAY_HIJACK_HOST` | (Optional) `false` to skip DNS hijacking on the host (VPN/adblocking for Tailscale peers only) |
| `GATEWAY_MSS` | (Optional) TCP MSS clamp value (default 1120); 1180 for speed on healthy paths |
| `WARP_FAIL_THRESHOLD` | (Optional) Consecutive `warp=off` polls before auto-shutdown (default 6 = 30s; 3 = 15s) |
| `WARP_ENDPOINT_1` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 1 (default `162.159.192.1`) |
| `WARP_ENDPOINT_2` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 2 (default `162.159.193.1`) |
| `KILL_SWITCH` | (Optional) `true` to enforce strict PF lock  drops all host traffic if VPN fails. Breaks SSH if VPN dies. |
| `HOST_LEAK_PROBE` | (Optional) `false` to disable the 300ms host-egress leak prober (default: `true`) |
| `STOP_COLIMA_ON_EXIT` | (Optional) `true` to fully shut down the Colima VM on toggle-off (saves battery on dedicated hosts) |
| `ADGUARD_USER` | (Optional) AdGuard Home web UI username (default: `admin`) |
| `ADGUARD_PASSWORD` | (Optional) Shared password for AdGuard Home and Tor ControlPort (default: `nullexit`) |
```

```
| `BLOCKED_COUNTRIES` | Space-separated 2-letter ISO codes to block via ipdeny.com CIDR ranges (e.g. `"kp il cn ru"`) |

---

## 8. Diagnostic Commands

### Container Health
```bash
docker ps --format '{{.Names}} {{.Status}}'
docker compose exec -T tailscale tailscale status
docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace
```

SOCKS5 Proxy
```bash
curl --socks5-hostname 127.0.0.1:1080 --max-time 5 -s https://www.cloudflare.com/cdn-cgi/trace
```

AdGuard DNS
```bash
dig +tcp @127.0.0.1 -p 5354 google.com +short +timeout=5
```

Firewall & Routing
```bash
# nftables FORWARD chain
docker compose exec -T warp iptables -L FORWARD -v --line-numbers
# Legacy FORWARD chain
docker compose exec -T warp iptables-legacy -L FORWARD -v --line-numbers
# IP rules + routing tables
docker compose exec -T warp ip rule show
docker compose exec -T warp ip route show table 199
docker compose exec -T warp ip route show table 200
```

Host macOS
```bash
tailscale status
brew services list | grep tailscale
networksetup -getdnsservers Wi-Fi
networksetup -getsocksfirewallproxy Wi-Fi
sysctl net.inet.ip.forwarding
```

Part II Design Decisions & Threat Model

9. Privacy Architecture & Threat Model

Layered Defense
The gateway stacks four layers targeting different attack surfaces:

- **Layer 1 ISP-Level Surveillance (WARP):** Your ISP sees only an encrypted WireGuard tunnel to a Cloudflare IP. In the US, ISPs can legally sell your browsing metadata and are subject to secret NSLs (National Security Letters). WARP removes their visibility entirely.
- **Layer 2 DNS Tracking (AdGuard Home):** DNS is the phone book of the internet. By default, queries go to your ISP's resolver, giving them a complete log of your traffic. AdGuard intercepts all DNS from mesh devices, blocks trackers, and resolves queries through the WARP tunnel.
- **Layer 3 Device Identity & IP Exposure (Tailscale):** On untrusted networks (coffee shops, public Wi-Fi), your device's IP is exposed. Tailscale creates an encrypted WireGuard mesh between your devices, routing all traffic to your gateway exit node.
- **Layer 4 Kernel-Level IP Blocking (ipset/iptables):** Sophisticated malware bypasses DNS entirely with hardcoded IPs. The `rule-compiler` fetches ~16,700 threat-intelligence IPs/CIDRs (Spamhaus, Feodo Tracker, Emerging Threats, CINS) and enforces them as `DROP` rules in the kernel `FORWARD` chain blocking both outbound C2 connections and inbound attack traffic.

The Documented Threat: Why This Matters
This architecture is calibrated against documented mass-surveillance programs revealed in the Snowden disclosures:
- **PRISM:** Bulk collection of communication content from major tech providers.
- **UPSTREAM:** Tapping physical backbone fiber, collecting metadata and content in bulk.
- **XKeyscore:** Retroactive search of unencrypted traffic.
- **MUSCULAR:** Infiltration of unencrypted internal datacenter interconnect links.

Passive bulk collection is not paranoia it is a documented reality. By double-encrypting and routing traffic, nullexit prevents your ISP from harvesting your metadata.

Trust Assumptions & Shifting
Every component shifts trust rather than eliminating it:
```

```
1. **Physical Device Integrity:** You trust your physical devices are not compromised.
2. **Cloudflare Integrity:** You trust Cloudflare not to correlate your WARP tunnel with exit traffic.
3. **Tailscale Integrity:** You trust Tailscale's coordination server not to MITM WireGuard keys.

The goal: no single provider has a complete picture. Your ISP sees an encrypted tunnel. WARP sees traffic with no account identity. Tailscale sees node topology but not content.

### 9.1 Exit Node IP Rotation (Design Decision)

A common question for privacy architectures is whether the system should aggressively rotate its public exit IP (e.g., every 5 minutes, like a Tor circuit or a proxy rotator).

`nullexit` uses Cloudflare WARP via Gluetun. WARP provides **anonymity in a crowd** rather than aggressive rotation. Thousands of users share the same Cloudflare Edge IP simultaneously, making your traffic indistinguishable from the herd. The IP may occasionally rotate on its own (which `host-leak-probe.sh` logs as a `ROTATE` event), but it remains generally stable.

**Why aggressive IP rotation was intentionally excluded:**
1. **Broken TCP Connections:** Every rotation instantly drops all long-lived TCP sessions (SSH, large downloads, WebSockets, gaming, Zoom).
2. **CAPTCHA & 2FA Hell:** Modern web security (e.g., Cloudflare, Akamai) aggressively flags IP-hopping as bot behavior. Aggressive rotation guarantees the user will be bombarded with "Verify you are human" CAPTCHAs and "New Login Detected" 2FA emails constantly.
3. **WireGuard Architecture:** WireGuard is stateless and does not natively support IP hopping on the fly. To rotate an IP, the VPN client must actively drop the tunnel, request a new endpoint, and re-handshake. This introduces latency gaps where the kill-switch would repeatedly trigger and blackhole the user's internet.

**Verdict:** For a daily-driver secure gateway, shared-IP crowding (WARP) is superior to aggressive rotation. As an added benefit, Cloudflare WARP IPs are actually highly static (tied to the persistent WireGuard public key generated for your device identity). Because your IP stays static and is part of Cloudflare's own trusted network, it builds a high "trust score," virtually eliminating CAPTCHA loops while still hiding you within the massive crowd of other users in that data center. All mesh devices routing through your gateway will reliably share this exact same trusted IP.

*(Note: The only exception where aggressive IP rotation is genuinely required is for specialized offensive security taskslike high-volume web scraping or dark web researchwhere avoiding IP bans or active tracking overrides the need for TCP stability.)*

### 9.2 Cloudflare WARP Privacy & Logging

Because `nullexit` uses Cloudflare WARP (via Gluetun) as its upstream exit node, the system inherits Cloudflare's consumer privacy policy.

**What Cloudflare DOES NOT log (Strict No-Logs Policy):**
* **Browsing History:** They do not log the domains or URLs you visit.
* **Traffic Destinations:** They do not log the IP addresses of the servers you connect to.
* **Content:** They do not inspect or log the payload/data of your traffic.
* **Data Brokering:** They explicitly guarantee they will never sell or rent your data to advertisers.

**What Cloudflare DOES log (Minimal Telemetry):**
* **Data Transfer Volume:** Total bandwidth used (required to enforce 10GB/mo quotas if you are not using a paid WARP+ key).
* **Performance Metrics:** Anonymized connection speeds and latency to optimize their network routing.
* **Aggregate Volume:** Total traffic volumes to popular destinations in aggregate (e.g., "datacenter X sent 50 TB to Netflix"), stripped of all user IDs.
* **Device ID (Public Key):** They store the persistent WireGuard public key generated by your container to authorize the connection.

**Privacy Verdict:**
Cloudflare WARP provides excellent **historical privacy**. Because they do not log your destinations, they cannot comply with historical subpoenas asking what websites you visited yesterday. However, it is **not absolute anonymity**. They still know your real ISP IP address while you are actively connected. For daily-driver privacy against ISPs, hackers, and corporate trackers, it is top-tier. For nation-state evasion, use Tor (see 11).

### 9.3 OPSEC: Python Runtime Airgap (Isolated Mode)
`nullexit` explicitly relies on the host's native system Python 3 runtime for critical components like `scripts/dns-proxy.py`. Because this script runs with elevated privileges (`sudo -n`) to bind to the privileged UDP/53 socket, modifying or polluting this Python environment is strictly forbidden.

**The Zero-Dependency Rule & Isolated Mode (`-I`)**
You must **NEVER** install third-party `pip` packages (e.g., `requests`, `pyyaml`) into the system Python environment that `nullexit` uses.
- All Python scripts in this repository must be written using *only* the Python Standard Library (`socket`, `urllib`, `json`, `ssl`).
- Third-party packages introduce a massive supply-chain attack vector. A compromised `pip` package installed globally could allow local malware to hook into the Python runtime and exploit the `sudo` privilege of the DNS proxy to gain full root access.
- To enforce strict immunity against user-installed malicious packages, `toggle.sh` explicitly invokes the proxy using **Python's Isolated Mode** (`python3 -I`). This flag forces the interpreter to completely ignore `PYTHONPATH`, `PYTHONHOME`, and the user's `site-packages` directory, executing exclusively from the
```

```
     read-only, Apple-signed system framework.

### 9.4 Recommended Upgrades

| Layer | Default | Upgraded |
|---|---|---|
| VPN Exit | Cloudflare WARP (US) | Mullvad (Swedish, proven no-logs, anonymous payment) |
| DNS Resolver | AdGuard via WARP | AdGuard via Mullvad DoH (Cloudflare-blind) |
| Coordination | Tailscale (US, OAuth) | Headscale (self-hosted, no third-party) |
| Auth | OAuth / persistent keys | Ephemeral auth keys (no identity persistence) |
| Content | WireGuard asymmetric | WireGuard + PSKs (coordination server blind) |

#### Mullvad (Replacing WARP)
Swedish-based. Police raided their offices in 2023 and left empty-handed  no logs exist to seize. Accounts are random numbers. Payment by cash or Monero. Replace `VPN_SERVICE_PROVIDER=custom` with `VPN_SERVICE_PROVIDER=mullvad` in `docker-compose.yml`.

#### Mullvad DoH Upstreams
Point AdGuard's upstream resolvers to `https://adblock.dns.mullvad.net/dns-query`. Cloudflare WARP only sees encrypted HTTPS to Mullvad's IPs  completely blind to your DNS queries.

#### Headscale
Self-hosted Tailscale coordination server. Eliminates Tailscale the company, keeping all node topology under your control.

#### Ephemeral Auth Keys
Generate in the Tailscale admin panel. Used once to authenticate; node disappears from the admin panel after disconnect. Breaks persistent identity linkage.

#### Pre-Shared Keys (PSKs)
Add a WireGuard PSK between peer nodes. Adds a symmetric encryption layer that Tailscale's coordination server has no knowledge of, blinding it from reading transit content.

### 9.5 Structural Limitations
- **Traffic Analysis / Metadata:** Passive fiber tapping (UPSTREAM) can observe timestamps, data volumes, and IP ranges without reading content. Defeating traffic analysis requires mixing networks (Tor) or continuous dummy packet padding  both heavily degrade performance.
- **Targeted State-Level Adversaries:** Consumer-grade VPNs are not a complete shield against a nation-state actively targeting you. This stack is designed to defeat bulk passive surveillance and corporate dragnet collection.

---

## 10. Per-Device Access Control

Because every mesh device routes internet-bound traffic through the `warp` container's `FORWARD` chain, they all appear with their stable `100.x.x.x` Tailscale IP as the `src` address. This gives two powerful enforcement surfaces:

### IP & Port Level (`scripts/routing-fix.sh`)
Write raw `iptables` rules targeting specific devices. The 30-second idempotent loop auto-survives Gluetun reconnects:

```bash
# Block a specific device from a target subnet
iptables -I FORWARD -s 100.x.x.x -d 203.0.113.0/24 -j DROP

# Restrict a device to only HTTP/HTTPS
iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 3000 -j DROP
iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 443 -j DROP

# Block a device from internet egress entirely (mesh only)
iptables -I FORWARD -s 100.x.x.x -j DROP

# Time-based (e.g., cut off 10 PM  7 AM)
iptables -I FORWARD -s 100.x.x.x -m time --timestart 22:00 --timestop 07:00 -j DROP
```

DNS Level (AdGuard REST API)
AdGuard Home supports per-client rules keyed by Tailscale IP:

```bash
curl -X POST http://localhost:80/control/clients/add \
  -H "Content-Type: application/json" \
  -d '{
    "name": "kids-ipad",
    "ids": ["100.x.x.x"],
    "blocked_services": ["youtube", "tiktok", "instagram"],
    "safesearch_enabled": true
  }'
```
```

```
428
429 ---
430
431 ## 11. Tor & OPSEC Architecture
432
433 ### 11.1 Tor-over-WARP Topology
434
435 When extreme anonymity is required, integrating the Tor network into `nullexit` provides a powerful Tor-over-
      VPN topology:
436 * **The University/ISP** sees only an encrypted Cloudflare WireGuard tunnel, rendering them completely blind to
       the fact that you are using Tor (effectively bypassing Enterprise firewall Tor blocks).
437 * **Cloudflare** sees an encrypted TCP stream heading to a Tor Guard Node. They cannot see your DNS lookups,
      the destination website, or the content of your traffic.
438 * **The Tor Exit Node** sees your traffic, but not your real IP (only the Cloudflare IP).
439
440 #### The "Transparent Proxy" Trap (What NOT to do)
441 It is tempting to build a system-wide "Transparent Tor Proxy" that forces all host traffic (e.g., standard
      Chrome/Safari browsers, background iCloud syncs) through Tor automatically. **This is a massive OPSEC trap
      and destroys anonymity.**
442 Modern operating systems and browsers will instantly leak your identity through background syncing (e.g., Apple
       Mail, Google Drive) and browser fingerprinting (canvas fingerprints, WebRTC, screen resolution). The Tor
      network protects your *IP*, but if your payload contains identifying cookies or unique fingerprints, the
      Tor Exit Node (which could be run by malicious actors) will instantly tie your session to your real
      identity.
443
444 #### The nullexit Implementation: Isolated Procedural Proxy
445 To safely integrate Tor without triggering the transparent proxy trap, `nullexit` implements an **Isolated Tor
      SOCKS5 Proxy** with **Procedurally Generated Ports**:
446 1. **Isolated Container:** A lightweight `tor` Docker container runs securely inside the `warp` network
      namespace. It does not hijack system traffic.
447 2. **Opt-In Usage:** You must consciously configure a highly-hardened browser (like the official Tor Browser or
       a dedicated Firefox profile) to use the proxy. This prevents accidental background sync leaks.
448 3. **Procedurally Generated Ports:** Rather than hardcoding the standard proxy ports (like `127.0.0.1:9050` or
      `1080`), `toggle.sh` randomizes all internal proxy ports on every boot into the ephemeral range
      (10000-65000) and silently writes them to `/tmp/nullexit-ports.env`. This completely defeats static
      fingerprinting by malware.
449 4. **Kernel-Level Collision Avoidance:** To prevent the randomizer from picking a port already in use by a root
       daemon or Apple service (which would crash the gateway), the script directly queries the macOS kernel
      socket table (`netstat -an -p tcp`) to verify the port is absolutely free before assigning it.
450 5. **The Honey-Port Tripwire:** While procedural ports defeat static fingerprinting, advanced malware might
      attempt a full sequential port scan of `localhost` to find the hidden proxy. To counter this, `toggle.sh`
      spawns a "Honey-Port" (a fake random port with a netcat listener attached). If any local application
      attempts to connect to the Honey-Port, it instantly logs a critical security alert to `output.log` and
      fires a macOS desktop notification, serving as an early-warning Intrusion Detection System (IDS) for local
      malware.
451
452 > **Threat Model note:** The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). Port
      randomization and the Honey-Port only defeat blind, generic port-scanners  they do NOT protect against a
      targeted local process that reads `/tmp/nullexit-ports.env`, greps process memory, or runs `lsof -i`. The
      proper mitigation is a host-level loopback-filtering firewall (LuLu/Little Snitch). See 15.9 (Threat Model:
       Local Proxy Discovery) for the full analysis and the rogue-binary verification test.
453
454 ### 11.2 Hardening: Tor Bridges and obfs4 Pluggable Transports
455 For users operating under severe Deep Packet Inspection (DPI), `nullexit` supports configuring Tor to connect
      exclusively through private bridges using the `obfs4` pluggable transport.
456
457 When `TOR_USE_BRIDGES=true` is set:
458 - **What it does:** It hides the entry IP from being matched against the public Tor consensus list, and
      disguises the traffic's protocol fingerprint from the WARP exit provider (Cloudflare).
459 - **What it does NOT do:** It does *not* add any extra layers of anonymity beyond what Tor already provides. It
       is purely a censorship-circumvention and obfuscation mechanism.
460 - **Limitations:** `obfs4` is highly effective against passive DPI, but it is not guaranteed to defeat active-
      probing-resistant detection by a highly sophisticated state-level adversary.
461
462 > **Architecture Rule:** The `nullexit` gateway deliberately does NOT bundle, hardcode, or automatically fetch
      bridge lines. Bridge lines are highly sensitive and expire. Users must obtain their own bridge lines from `
      https://bridges.torproject.org` and manually populate the `tor-bridges.txt` file. If bridges are enabled
      but the file is missing or empty, the Tor container will intentionally crash and fail to start rather than
      silently falling back to public guard relays.
463
464 **The Bootstrapping Paradox (Out-of-Band Key Exchange)**
465 Automating the bridge acquisition process (e.g., scripting the gateway to log into the Telegram `@GetBridgesBot
      ` via MTProto API) introduces catastrophic OPSEC flaws:
466 1. **Identity Leaks:** Baking a Telegram session into the container explicitly links the "anonymous" gateway to
       a real-world physical phone number.
467 2. **The Chicken & Egg Deadlock:** In heavily censored regimes, Telegram itself is usually blocked. If the
      container requires Telegram to fetch bridges to bypass the firewall, it will fail because it cannot bypass
      the firewall to reach Telegram.
468 3. **Bridge Exhaustion & CAPTCHAs:** Automated scraping behaves identically to state-sponsored scrapers,
      triggering Tor's anti-bot CAPTCHAs which leads to silent proxy failures and exhausts the limited public
      bridge pool.
```

```
469
470 Instead, users should rely on a less secure layer to manually reach Telegram, obtain the bridges, and populate
      `tor-bridges.txt`. This can be done via the standard `nullexit` WARP tunnel, a mobile phone on cellular
      data, Mullvad VPN, or a secure remote VPS. *(Note: We plan on fully integrating Mullvad and VPS
      architectures directly into `nullexit` in the future, but have not had the time to build it yet).* This
      intentionally "air-gaps" the cryptographic key exchange from the proxy infrastructure, leaving zero API
      tokens and zero network traces for an adversary to exploit.
471
472 ### 11.3 Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification
473
474 #### Context
475 Unlike the host's SOCKS5 fallback path (whose kill-switch interaction is documented in 15.7), the `tor`
      container's connection status and fail-closed behavior under VPN tunnel drops is not governed by macOS `pf`
       rules. Instead, it relies on network namespace inheritance inside Docker Compose.
476
477 #### Mechanics
478 1. **Shared Network Namespace:** The `tor` service is configured with `network_mode: service:warp` in `docker-
      compose.yml`. This shares the network namespace of the `warp` (Gluetun) container.
479 2. **Inherited Egress Rules:** All socket communication within the namespace is bound to the same networking
      stack. Gluetun manages this stack by redirecting traffic through `tun0` (the WireGuard/WARP tunnel) and
      applying `iptables` rules that explicitly drop outbound traffic originating on `eth0` (the physical/bridge
      interface), except for permitted local subnets or the VPN server's endpoint.
480 3. **Tunnel Fail-Closed:** If the WARP connection drops or the `tun0` interface is brought down, the routing
      table entries for `tun0` become invalid or disappear. Because the `iptables` rules in the shared namespace
      continue to drop any outgoing internet traffic attempting to exit via `eth0`, the `tor` container cannot
      leak raw traffic to the host's underlying networks. Egress fails closed.
481
482 #### Verification / Manual Test Case
483 To manually verify that the Tor container fails closed and does not leak traffic when the tunnel dies:
484
485 1. **Verify active path:** Ensure the gateway is active (`toggle.sh start`) and fetch the Tor port (`
      $TOR_SOCKS_PORT` in `/tmp/nullexit-ports.env`). Verify Tor traffic routes correctly:
486    ```bash
487    curl --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip
488    ```
489    *(This should output a Tor exit node IP).*
490
491 2. **Simulate tunnel drop:** Bring down the virtual interface inside the shared network namespace:
492    ```bash
493    docker compose exec warp ip link set dev tun0 down
494    ```
495
496 3. **Re-run the egress check:**
497    ```bash
498    curl --socks5-hostname 127.0.0.1:$TOR_SOCKS_PORT https://check.torproject.org/api/ip
499    ```
500    * **Expected Result:** The request must hang and eventually time out, or immediately fail with a proxy error
         (e.g., `curl: (97) SOCKS5: connection failed` or `curl: (7) Failed to connect`).
501    * **Leaked Egress Check:** Check the host's Wi-Fi router or firewall logs. No outbound packets from the Tor
       container should route directly over the bridge network to the internet.
502
503 ### 11.4 Transparent .onion Routing via RFC 2544 Subnet (198.18.0.0/15)
504
505 #### Context
506 To enable seamless dark web browsing across the entire mesh, the gateway requires a mechanism to dynamically
      map `.onion` domains to IP addresses that the host operating system and network layers can route.
507
508 #### IP Address Conflict & Solution
509 1. **Initial Conflict:** Initially, the Tor virtual address network was configured to map `.onion` sites within
       the `10.192.0.0/10` range. However, the Colima Docker daemon dynamically allocated `10.200.1.0/24` to the
      containers. Because this Docker subnet mathematically falls within the `10.192.0.0/10` block, a routing
      loop occurred: all host packets destined for the Docker containers on ports like `9050` or `9051` were
      intercepted by the transparent proxy NAT rules and dropped, causing proxy failures.
510 2. **Tailscale Private IP Exclusion:** Shifting the subnet to another private range like `10.123.0.0/16`
      resolved the port conflict, but led to connection timeouts. On macOS, Tailscale's Exit Node routing
      actively excludes RFC 1918 private subnets (e.g. `10.0.0.0/8`, `172.16.0.0/12`, `192.168.0.0/16`) to
      maintain accessibility to local LAN resources (like printers/routers). Thus, packets targeting `10.123.x.x`
       were dropped locally by Tailscale before reaching the gateway tunnel.
511 3. **The Benchmarking Subnet Solution:** To bypass these private IP exclusions without manual static routing
      tables, Tor's virtual network was changed to **`198.18.0.0/15`** (reserved by RFC 2544 for benchmarking).
      Since this is not a private RFC 1918 range, Tailscale exit-node routing automatically encapsulates and
      routes all traffic destined for `198.18.0.0/15` through the tunnel to the gateway.
512 4. **Transparent NAT Redirection:** The `routing-fix` sidecar applies NAT rules in the shared network namespace
      :
513    ```bash
514    iptables -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040
515    ```
516    This intercepts all incoming TCP traffic targeting the mapped `.onion` range and transparently proxies it
      through Tor's transparent port (`9040`).
517
518 #### Exit Node Exclusions
```

The gateway supports the ability to exclude specific countries from being used as exit nodes (e.g. to avoid certain jurisdictions or nodes with poor network latency). This is configured globally via the `TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`, which gets dynamically appended to Tor's `ExcludeExitNodes` and strictly enforced with `StrictNodes 1`.

---

12. Censorship-Resistant Transport & Egress Compartmentalization

12.1 Why this is needed
WARP can be blocked from deployment countries with strict internet controls. The most likely cause isn't that "encrypted traffic is blocked" in general it's that WireGuard has a recognizable handshake signature, and `WIREGUARD_ENDPOINT_IP=162.159.192.1` on `WIREGUARD_ENDPOINT_PORT=2408` is a fixed, easily-enumerable target (Cloudflare's known WARP range). That combination is exactly what IP/ASN-based and DPI-based blocking is good at catching.

12.2 Important correction: Gluetun's `SHADOWSOCKS=on` is the wrong direction
Gluetun's built-in Shadowsocks option runs a Shadowsocks **server** inside the already-connected VPN tunnel, so other LAN devices can reach out through Gluetun's *existing* connection via the SS protocol. It is not a way to make Gluetun itself connect **outbound** through a Shadowsocks server Gluetun only speaks OpenVPN or WireGuard as its actual transport. There is no `VPN_TYPE=shadowsocks`. Confirmed against Gluetun's own maintainer/community discussion: people have asked for a "Shadowsocks-client" mode specifically to chain through Gluetun, and the standing answer is "run a separate Shadowsocks client container."

12.3 Diagnose before building anything
The right fix depends entirely on what's actually being blocked:
- Point Gluetun at a **different provider/IP/ASN** (any other Gluetun-supported WireGuard/OpenVPN provider, or a self-hosted WireGuard server) and test. If that gets through, the block is Cloudflare/WARP-specific, not a general WireGuard block **no new component needed**, just swap `VPN_SERVICE_PROVIDER` / the endpoint.
- If WireGuard fails regardless of endpoint, the block is protocol-level (DPI fingerprinting the handshake) proceed to obfuscation.

12.4 Path A: Swap WARP for a different Gluetun-native transport
Simplest fix, only applies if 12.3 shows the block is Cloudflare-specific. Change `VPN_SERVICE_PROVIDER` / `VPN_TYPE` / endpoint in `docker-compose.yml` and in the now-centralized `WARP_ENDPOINT_*` values in `.env`. Nothing else in the stack needs to change.

12.5 Path B: Add an obfuscation hop
Needed if WireGuard itself is being fingerprinted/blocked. Two ways to slot it in:

B1 Wrap (recommended): Run a Shadowsocks (or obfs4/v2ray) client as a new sidecar container. Point Gluetun's `WIREGUARD_ENDPOINT_IP` at `127.0.0.1:<local-port>` instead of Cloudflare directly; the sidecar relays that traffic to a Shadowsocks server you run yourself abroad. Keeps Gluetun's kill-switch, healthchecks, and the entire rest of the stack (Tailscale, AdGuard, ipset, `routing-fix.sh`) untouched, since none of them know the transport changed underneath `warp`.

B2 Replace: Swap the `warp` service entirely for a Shadowsocks-client + `tun2socks` container that owns the network namespace instead of Gluetun. More invasive loses Gluetun's built-in kill-switch/healthcheck, which would need to be rebuilt by hand (see Section 12.9 for the pluggable architecture to support this natively).

Recommendation: B1. Smaller surface area, preserves the self-healing architecture already built and documented.

12.6 Fingerprinting caveat
Plain Shadowsocks can still be caught by active-probing DPI in more aggressive censorship environments. If plain SS gets blocked too, the next step up is an obfuscation plugin (`v2ray-plugin`, `Cloak`) or a newer protocol designed against active probing (VLESS+Reality, Hysteria2). Test plain SS first don't over-build before confirming it's needed.

12.7 AmneziaWG (The High-Speed Alternative)
If you want to maintain the bare-metal speeds of WireGuard while bypassing DPI (e.g. in Egypt or Russia), the best alternative to Shadowsocks/V2Ray is AmneziaWG. AmneziaWG is a fork of WireGuard that pads the handshake packets with randomized "junk" data, entirely destroying the 148-byte DPI signature that firewalls look for.
- **Pros:** Because it runs entirely as UDP without TCP-over-TCP overhead, it completely avoids "TCP meltdown" latency spikes associated with Shadowsocks/V2Ray wrappers.
- **Cons:** You cannot use Cloudflare WARP. Furthermore, it requires completely ripping Gluetun out of the `nullexit` stack and building a custom Docker container, meaning you lose Gluetun's built-in kill-switches and health-check logic.

12.8 Infrastructure Costs & Privacy
To use either Shadowsocks (Path B) or AmneziaWG, you cannot use the free Cloudflare WARP infrastructure, because Cloudflare servers only speak standard WireGuard. The software for both custom protocols is 100% free and open-source, but you must host the remote server infrastructure yourself.
- **Free Tier (Oracle Cloud):** You can host your remote server on Oracle Cloud's "Always Free" tier for $0. However, this routes your highly sensitive, censorship-evading traffic directly through Oraclea massive corporate data broker. If privacy is the core objective, handing all your metadata to Oracle defeats the purpose.
- **Paid Tier (Hetzner / Mullvad):** The recommended path is to rent a standard ~$4/month Virtual Private Server (VPS) in a free country (e.g., Hetzner in Germany, subject to strict EU GDPR privacy laws) and self-

```
     host the server. Alternatively, Mullvad VPN (5/month) provides native v2ray/Shadowsocks bridges built into
     their network, allowing you to bypass censorship with a strict zero-logs policy. It is highly recommended
     to pay with untrackable payment methods to maintain complete anonymity, which these providers typically
     support without requiring local residency.
560
561  ### 12.9 The Future: Egress Compartmentalization (Pluggable Architecture)
562  To natively support invasive replacement paths like **Path B2** or **AmneziaWG** (Section 12.7) without
     breaking the core `nullexit` routing and monitoring logic, the architecture must be refactored into a
     modular, "pluggable" design.
563
564  1. **The Egress Plugin System:** Currently, WARP-specific logic (like `cdn-cgi/trace` health checks and
     hardcoded interface names) is scattered across `toggle.sh`, `scripts/routing-fix.sh`, and `scripts/host-
     leak-probe.sh`. This should be abstracted into a `scripts/plugins/` directory, where each egress method (`
     warp.sh`, `v2ray.sh`) implements standardized functions: `get_egress_interface()`, `check_health()`, and `
     get_bypass_ips()`.
565  2. **Dynamic Docker Compose:** Based on an `EGRESS_TYPE` variable in `.env`, dynamic compose file generation
     would spin up only the required egress containers (e.g., skipping the `warp` container when `EGRESS_TYPE=
     v2ray` is set).
566  3. **Agnostic Routing & Watchers:** `routing-fix.sh` would dynamically read the `EGRESS_INTERFACE` from the
     active plugin to apply `iptables` rules, and the background watchers in `toggle.sh` would become an
     agnostic `Egress Watcher` that invokes `check_health()` periodically.
567
568  This compartmentalization transforms `nullexit` into a universal, censorship-resistant gateway where the entire
      egress layer can be swapped by changing a single `.env` variable and running `./toggle.sh --restart`.
569
570  ---
571
572  ## 13. Roaming & Recovery Design: How nullexit Mimics Commercial VPNs
573
574  The `nullexit` roaming/network recovery subsystem replicates the exact engineering mechanics used by commercial
     , premium VPN clients (like Mullvad or NordVPN) to transition across Wi-Fi networks safely and seamlessly.
     This is the design overview; the specific bugs encountered and fixed along the way live in the Incident Log
      (15.5 Roaming/Sleep-Wake, 15.7 Tailscale P2P/SNAT).
575
576  ### The 4-Step Transition Cycle
577  When your Mac disconnects from one Wi-Fi and associates with another, `nullexit` coordinates the following
     events:
578
579  1. **Firewall Lock (Kill-Switch)**: The macOS Packet Filter (PF) anchor `com.apple/nullexit` remains locked on
     the physical interface (`en0`), blocking all direct outbound IPv4 and IPv6 traffic. This guarantees that
     your real ISP IP or unencrypted DNS requests **never leak** onto the new network while the connection is
     rebuilding.
580  2. **System Event Interception**: A launchd LaunchAgent (`watcher.sh`) listens to the macOS System
     Configuration framework for the `State:/Network/Global/IPv4` key changes, spawning `recover.sh --post-wake`
      in under 500ms when a change is detected.
581  3. **Bypass Routing**: The script dynamically resolves the new physical network's IP gateway (e.g.
     `192.168.137.1`), cleans up stale bypass routes from the previous network, and adds static bypass routes
     pointing to Cloudflare's WARP endpoints and the Tailscale control plane directly via the new gateway. This
     allows the VPN client to negotiate its handshake outside of the tunnel.
582  4. **Hole-Punching / NAT Re-negotiation**: The script runs a target check on the `warp` container's tunnel. If
     the WireGuard connection is wedged by the network switch, it force-recreates the container to establish a
     fresh NAT binding to Cloudflare, restoring connection in ~15 seconds.
583
584  ---
585
586  ## 14. Deployment, Packaging & Known Limitations
587
588  ### 14.1 Packaging nullexit as a macOS Application (.dmg)
589
590  nullexit can be packaged into a native `.app` bundle, but this introduces nested virtualization requirements.
591
592  nullexit uses **Colima** (Apple Virtualization Framework `vz`) to run a Linux VM hosting Docker. Installing the
      `.app` inside a virtualized macOS environment (Parallels, cloud Mac) means running a Linux VM inside a
     macOS VM  requiring hardware-level **nested virtualization**.
593
594  #### Hardware Support
595  - **Supported:** Apple Silicon M3/M4 (macOS Sequoia 15+), Intel Macs with VMX passthrough.
596  - **Unsupported:** M1, M2, A18 Pro. Colima crashes silently; the terminal (`setup.sh`) surfaces crash logs.
597
598  #### Build Steps
599  ```bash
600  # 1. Verify nested virtualization support
601  sysctl -a | grep hv_nested_virt_supported  # expect: 1
602
603  # 2. Compile the AppleScript launcher
604  osacompile -o "Nullexit.app" "Toggle Gateway.applescript"
605
606  # 3. Embed scripts into app resources
607  mkdir -p "Nullexit.app/Contents/Resources/scripts"
608  cp toggle.sh setup.sh recover.sh docker-compose.yml "Nullexit.app/Contents/Resources/"
609
```

```
# 4. Package into DMG
hdiutil create -volname "Nullexit Installer" -srcfolder "./Nullexit.app" -ov -format UDZO Nullexit.dmg

# 5. Bypass Gatekeeper (unsigned app)
xattr -cr /Applications/Nullexit.app
```

Without a paid Apple Developer certificate ($99/yr), users see an "App is damaged" Gatekeeper warning and need to run step 5.

14.2 Moonlight/Sunshine + Tailscale Race Condition

When streaming from a remote Windows host running Sunshine over a Tailscale mesh (e.g., while the client is on a cellular hotspot), a race condition can occur on boot:

1. Both Tailscale and Sunshine are set to start automatically on boot.
2. Tailscale takes a few seconds to fully initialize its `100.x.x.x` virtual adapter and establish a connection.
3. **Sunshine starts too fast.** It launches before Tailscale is fully ready. Since Sunshine only binds to network interfaces that are active at the exact moment of startup, it misses the Tailscale adapter entirely.
4. Moonlight clients (even when configured with the correct Tailscale IP) will fail to connect because Sunshine isn't listening on that interface.

The "Magic Toggle" Symptom:
If the user manually enables or disables Wi-Fi on the host PC, it triggers a global Windows network change event. Sunshine listens for these events, re-enumerates all network adapters, and says, "Oh, there's a Tailscale adapter here!" and finally binds to it. Suddenly, Moonlight can connect.

The Permanent Fix:
On the Windows host, open `services.msc`, locate the **Sunshine Service**, and change its Startup type from **Automatic** to **Automatic (Delayed Start)**. This forces Windows to wait a minute or two after booting before launching Sunshine, ensuring Tailscale's virtual adapter is fully initialized first.

14.3 Chrome Remote Desktop Performance (UDP NAT)

Observation (July 11, 2026)
When the user connects to their host Mac using Chrome Remote Desktop (CRD) while `nullexit` is active, the connection suffers from massive latency, lag, and degraded video quality. AdGuard Home's `blocked.log` shows zero blocked DNS requests for Google, WebRTC, or `chromoting` endpoints, indicating this is not a DNS sinkhole issue. *(See also 15.11 for the separate CRD-on-remote-device connection-failure quirk.)*

Root Cause
This is a fundamental limitation of tunneling all host traffic through a strict commercial NAT (Cloudflare WARP).
1. Chrome Remote Desktop attempts to use **STUN (UDP hole-punching)** to establish a lightning-fast, direct peer-to-peer WebRTC connection between the client device and the host Mac.
2. Because all non-Tailscale traffic is forcefully routed through the `warp` container, the STUN packets exit the network from a Cloudflare IP.
3. Cloudflare's enterprise NAT infrastructure strictly drops unsolicited inbound UDP packets. The UDP hole-punch fails.
4. Because a direct P2P connection cannot be established, CRD falls back to **Relay Mode** (TURN), bouncing all video data back and forth through Google's cloud servers instead of sending it directly over the local or mesh network. This relaying introduces massive latency.

Workaround / Fix
This lag is the accepted "tax" for maintaining absolute cryptographic anonymity; the only way to restore CRD speed would be to leak its UDP traffic outside the tunnel (violating zero-trust).
To achieve low-latency remote access, users should bypass CRD entirely and use macOS's built-in **Screen Sharing (VNC)** directly over the Tailscale IP. Tailscale's sophisticated custom DERP/WireGuard architecture successfully UDP hole-punches through strict NATs, providing a direct, encrypted, high-speed connection without relying on external cloud relays.

14.4 Dynamic IP Fetch vs. MagicDNS

Observation
The `nullexit` gateway uses a dynamically fetched IP address (cached in `.gateway_ip`) to configure host routing and the `pf` kill-switch, rather than utilizing Tailscale's user-friendly MagicDNS hostname (e.g., `nullexit-gateway`).

Why MagicDNS is Not Used
1. **The "Chicken and Egg" Bootstrapping Problem:**
   macOS's Packet Filter (`pf`) and low-level routing commands (`route add`) operate strictly at Layer 3 and require raw IP addresses. If `toggle.sh` or `recover.sh` attempted to use `nullexit-gateway`, the host Mac would need to perform a DNS lookup to resolve it. However, the very first step of the gateway's initialization is to completely hijack the host's DNS and point it *at the gateway itself*. The Mac cannot resolve the gateway's MagicDNS name if it needs the gateway's IP to know where to send the DNS query!
2. **Dynamic Self-Healing (Avoiding Stale Cache):**
   To solve the bootstrapping problem without causing bugs if the node is deleted and re-authenticated on the Tailscale Admin Console, `toggle.sh` explicitly runs `docker compose exec ... tailscale ip -4` during boot to fetch the absolute freshest IP dynamically. It saves this to `.gateway_ip` so that fast-roaming scripts

```
     like `recover.sh` can instantly retrieve the routing target without incurring the latency of querying
     Docker.
660
661  ### 14.5 Battery & Power Measurement Quirks
662
663  When trying to optimize this framework (or any daemon) for battery life, developers often look for a way to
     programmatically measure the exact Wattage consumed by a specific process (e.g., "How many Watts is `toggle
     .sh` using?").
664
665  **This is a hardware impossibility on macOS and Linux.**
666
667  #### Why Activity Monitor is Lying to You
668  Your machine's logic board only has physical power sensors for the *entire* CPU package, the GPU, and the RAM.
     It knows the CPU is pulling 5W total, but it has no physical hardware capability to know whether Docker is
     using 3W and Chrome is using 2W.
669
670  The "Energy Impact" score you see in macOS Activity Monitor is a **synthetic, heuristic score** calculated by `
     powerd`. It penalizes apps for waking up the CPU (idle wakes), doing disk I/O, or keeping the screen awake.
      It is a relative score, not actual Wattage.
671
672  #### The Proper A/B Testing Methodology
673  If you want to truly know how many Watt-hours or mAh your background daemons (`routing-fix.sh`, `host-leak-
     probe.sh`, etc.) are costing you, you must perform a hardware-level A/B drain test:
674
675  1. Unplug the laptop, run the gateway, and note your exact battery mAh using:
676     ```bash
677     ioreg -l | grep "AppleRawCurrentCapacity"
678     ```
679  2. Leave the laptop idle for exactly 1 hour.
680  3. Run the command again to see exactly how many mAh were drained.
681  4. Turn the gateway off and repeat the test for another hour.
682
683  The delta in mAh drained between the two tests is the true, exact cost of running the software. When optimizing
      polling loops (e.g., changing `sleep 5` to `sleep 30`), this is the only reliable way to measure the
     actual battery life returned to the user.
684
685  ### 14.6 Host Lockdown Mode (Why It's Impossible on macOS)
686
687  A common privacy feature request is a "Host Lockdown Mode" (Mode 3): A mode where the Docker exit node remains
     perfectly functional for remote devices, but the host machine itself is completely blocked from accessing
     the internet (to prevent even encrypted host traffic from leaking metadata).
688
689  **Why this is impossible on macOS:**
690  On Linux (e.g., a Raspberry Pi), Docker runs natively. The host's traffic traverses the `OUTPUT` iptables chain
     , while Docker traffic traverses the `FORWARD` chain. We could easily drop all `OUTPUT` traffic to kill the
      host's internet, and Docker would continue routing traffic perfectly.
691
692  However, Docker on macOS runs inside a Linux Virtual Machine (via `qemu` or Apple `Virtualization.framework`).
     This VM runs as a standard macOS user-space application. If you configure the macOS firewall (`pf`) to drop
      all host outbound traffic, it will indiscriminately drop the VM application's traffic as well, instantly
     killing the Cloudflare WARP tunnel and the exit node.
693
694  Writing complex macOS `pf` rules to "allow the VM app, allow Cloudflare WARP IPs, allow Tailscale DERP IPs, but
      block everything else" is highly fragile and heavily discouraged. Since the current `toggle.sh` default
     mode already fully encrypts the host's traffic via the WARP tunnel, the host is already protected from ISP
     leaks.
695
696  ---
697
698  # Part III  Incident Log & Resolved Issues
699
700  This is the single, deduplicated log of every incident, bug, and resolved issue. Entries are organized into
     subsystem groups (15.115.12). Each entry follows a **Symptom / Root Cause / Fix** shape where applicable;
     packet-level analyses are preserved as clearly-labeled **Deep dive** blocks. The terse dated changelog
     lives in 17; this section holds the full post-mortems.
701
702  ## 15. Incident Log
703
704  ### 15.1 Exit-Node Return Path & SOCKS5 Failover
705
706  #### 15.1.1 Exit Node Return Path (`rx 0`) & The SOCKS5 Failover  RESOLVED
707  **Symptom:** Gateway showed `tx 936 rx 0`  transmitted but never received return traffic. For days, Tailscale's
      exit node refused to return traffic back to the client, leading to the assumption that Tailscale's
     userspace forwarding was bypassing `iptables` and ignoring the WARP `tun0` interface. In a desperate
     attempt to fix this, a **SOCKS5 proxy** (`scripts/socks5-proxy.py`) was introduced as a replacement.
708
709  **Root Cause:** Tailscale's userspace forwarding was *not* the problem. `docker-compose.yml` included `
     FIREWALL_OUTBOUND_SUBNETS=100.64.0.0/10`. Gluetun created a strict routing rule (`ip rule add to
     100.64.0.0/10 lookup 199`, priority 99) that forced all Tailscale CGNAT traffic to bypass the VPN via the
     unencrypted Docker bridge (`eth0`). Return packets were sent to the macOS host instead of back through `
     tailscale0`, blackholing all return packets back to the client.
```

```
**Fix:** Removed `FIREWALL_OUTBOUND_SUBNETS` entirely, immediately restoring the Tailscale exit node. `scripts/routing-fix.sh` now injects `lookup 52`, and return packets correctly flow back into `tailscale0`. Instead of removing the SOCKS5 proxy, we repurposed it into a **bulletproof failover** for when restrictive Wi-Fi networks block Tailscale UDP ports.

Architecture after the fix:
```
PRIMARY (Exit Node):
DNS/TCP/UDP: Apps Tailscale Exit Node gateway container tun0 WARP internet

FAILOVER (SOCKS5 - if Tailscale is blocked by local Wi-Fi):
DNS: Browser 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP
TCP: Apps SOCKS5:1080 gateway container tun0 WARP internet
Mesh: SSH/ping 100.x.x.x utun5 WireGuard peers (independent of proxy)
```

**Deep dive: Exit Node Return Path Analysis (June 26, 2026).**
This preserves the detailed packet-level analysis that led to identifying and fixing the `rx 0` bug.

*The exit node was probably never going to work in this container topology.* Here's the packet path traced by reading `ip rule show` and the routing tables live:

**Outbound (Phone-1  internet):**
1. Phone-1 sends packet to gateway via Tailscale mesh
2. Gateway's tailscale0 (kernel TUN, `TS_USERSPACE=false`) decrypts  packet enters FORWARD chain
3. Packet has src=`100.87.42.87` (Phone-1), dst=some internet IP
4. FORWARD chain (nftables): Rule 1 matches (`-i tailscale0 -j ACCEPT`)
5. Routing decision for the forwarded packet. Which ip rule matches?
   - Rule 98: `to 172.18.0.0/16`  no (dst is internet)
   - Rule 99: `to 100.64.0.0/10`  no (dst is internet)
   - Rule 100: `from 172.18.0.2`  **NO** (src is `100.87.42.87`, not the container IP!)
   - Rule 101: `not fwmark 0xca6c lookup 51820`  **YES** (no fwmark on forwarded packets)
6. Table 51820: `default dev tun0`  packet goes out through WARP
7. NAT POSTROUTING: `MASQUERADE on tun0`  src becomes WARP IP
8. Packet exits through WARP to internet  (in theory)

**Return (internet  Phone-1):**
1. Return packet arrives on tun0, conntrack de-MASQUERADEs dst back to `100.87.42.87`
2. Routing decision for dst=`100.87.42.87`:
   - Rule 99: `to 100.64.0.0/10 lookup 199`  **YES**
3. Table 199: `100.64.0.0/10 via 172.18.0.1 dev eth0`  **sends it to the Docker gateway!**
4. **This is wrong.** The packet should go to tailscale0 (to be re-encrypted and sent to Phone-1), but table 199 sends it out eth0 to the Docker bridge host.

**Table 199 is the smoking gun.** It has one route: `100.64.0.0/10 via 172.18.0.1 dev eth0`. This sends ALL Tailscale CGNAT return traffic out the Docker bridge to the host  not back through tailscale0. The return packet leaves the container's network namespace entirely, hits the host's routing stack, and gets dropped because the host has no idea what to do with a raw packet for `100.87.42.87`.

*Why table 199 exists and why it's wrong for exit node traffic.* Table 199 + the CGNAT rule (pref 99) was designed so that the **container itself** can reach other Tailscale peers (e.g., for mesh management, DERP relay). For that use case, routing CGNAT traffic via the Docker bridge to the host (which has its own tailscaled) makes sense. But for **forwarded exit node traffic**, the return packet needs to go back through tailscale0 inside the container  not out to the host. The CGNAT rule intercepts the return packet before it can be routed back to tailscale0 through the main table or Tailscale's own routing (table 52, which is at pref 5270).

*The SOCKS5 path works because it dodges all of this.* SOCKS5 proxy creates connections **from the container's own IP** (`172.18.0.2`). So:
- ip rule 100 (`from 172.18.0.2 lookup 200`) matches
- Table 200 has `default dev tun0`
- Return traffic comes back to `172.18.0.2` on tun0, conntrack de-NATs, proxy sends response to client

No FORWARD chain involved. No CGNAT rule involved. No table 199. It just works.

*Summary of findings:*

| Finding | Status | Evidence |
|---|---|---|
| Table 199 CGNAT route hijacks exit node return traffic | **Fixed** | Changed `lookup 199` to `lookup 52` in `scripts/routing-fix.sh`. Return packets now route via tailscale0. |
| FORWARD RELATED,ESTABLISHED missing | **Fixed** | Installed `iptables` via apk in `docker-compose.yml`. Rule now injects successfully. |
| DOCKER_GW variable parsing error | **Fixed** | Added `head -1` in `scripts/routing-fix.sh`. |
| `FIREWALL_OUTBOUND_SUBNETS` was the root cause | **Fixed** | Removed from `docker-compose.yml`. Gluetun no longer hijacks CGNAT routing. |
| Host tailscaled error state from aggressive `tailscale down` | **Mitigated** | Toggle script now detects and auto-restarts. |

#### 15.1.2 Container Default Route Bypasses WARP
```

```
**Problem:** Inside the WARP container's network namespace, the main routing table's default route was via `eth0` (Docker bridge), not `tun0` (WARP tunnel). While gluetun uses policy routing (rule 101  table 51820 `tun0`) for most traffic, traffic originating from the container's own IP (`172.18.0.2`) hits rule 100 (`from 172.18.0.2 lookup 200`), whose default was also via `eth0`. This caused the SOCKS5 proxy's outbound connections to bypass WARP.

**Fix:** Updated the `routing-fix` container to:
1. Change the main table's default route to `default dev tun0`
2. Change table 200's default route to `default dev tun0`
3. Add a specific route for the WARP WireGuard endpoint (`162.159.192.1`) through `eth0` via the Docker gateway in table 200, preventing a tunnel loop
4. Re-assert all routes every 30 seconds (gluetun may reset them on health checks)
5. Dynamically detect the Docker subnet from `ip route show dev eth0` instead of hardcoding `172.18.0.0/16`

#### 15.1.3 Hardcoded Docker Subnet in Routing Fix
**Problem:** The `routing-fix` container hardcoded `172.18.0.0/16` and `172.18.0.1` for the Docker bridge subnet and gateway. If Docker's bridge network changed (after `docker network prune`, Colima restart, or on a different machine), these routes would be wrong.

**Fix:** Detection is now dynamic using `ip route show`:
- `DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')`
- `DOCKER_GW=$(ip route show default | awk '{print $3}')`

This extracts the actual subnet and gateway directly from the routing table, working with any subnet size (`/16`, `/24`, `/20`, etc.).

### 15.2 Routing Loops & Subnet Collisions

#### 15.2.1 The Infinite Recursive Routing Loop (Hotspot Paradox)
**Problem:** A critical network topology crash occurs if the host Mac connects to a Mobile Hotspot (e.g., a Windows PC or Phone) that is *also* connected to the Tailscale mesh and has "Use Exit Node" enabled. Because the Mac is hosting the Exit Node, it routes its internet traffic to the Hotspot. The Hotspot intercepts the Mac's traffic on its way to the cellular tower and routes it *back* to the Exit Node (the Mac) via Tailscale. The Mac receives it, tries to send it out again, and the Hotspot intercepts it again. This creates an infinite, recursive encryption loop that instantly destroys network bandwidth and drops all connections.

**Fix:** This is a physical routing paradox that cannot be resolved in software. The upstream physical router providing internet to the Mac **must** be allowed to route its traffic directly to the physical WAN interface. The user must manually disable "Use Exit Node" on the upstream device providing the hotspot.

**Advanced Workaround:** If the upstream hotspot is a Windows machine and the user explicitly wants it to remain connected to the Exit Node, a permanent static route can be injected into the Windows routing kernel to forcefully bypass the Tailscale WFP interception for WARP packets. By binding the route to the physical adapter's Interface Index (IF), it becomes a permanent, roaming-aware bypass.

On Windows (Admin Command Prompt):
```cmd
route print  (find Interface List, note the IF number for the active Wi-Fi adapter, e.g. 15)
route -p add 162.159.192.1 mask 255.255.255.255 0.0.0.0 IF 15
route -p add 162.159.193.1 mask 255.255.255.255 0.0.0.0 IF 15
```

On Linux (Root Terminal):
```bash
ip route add 162.159.192.1 dev wlan0
ip route add 162.159.193.1 dev wlan0
```

On macOS (Root Terminal):
```bash
route add -host 162.159.192.1 -interface en0
route add -host 162.159.193.1 -interface en0
```

(Note: If the upstream router is an iPhone or Android device, you cannot inject static routes without jailbreak/root. You must simply disable "Use Exit Node" in the mobile Tailscale app).

This punches a tiny hole straight through the paradox, letting the Mac's WARP packets escape to the physical internet while keeping the rest of the upstream router's traffic securely trapped in the Tailscale Exit Node.

15.2.2 Infinite Egress Routing Loops & Subnet Takeover Paradoxes
Problem: Two distinct network loops can break host egress connectivity entirely when the exit node is active:

1. **The Recursive Host-VM Tunnel Loop:**
   - **Mechanism:** The host Mac's default route points to the Tailscale exit node (`utun*`). The exit node container sends its encrypted tunnel packets (destined for the WARP endpoint `162.159.192.1`) back to the host via the VM NAT. Without a static route exception on the host Mac, the host Mac routes those packets right back into the exit node (`utun*`), creating an infinite loop that blackouts all host network egress.
```

```
827    - **ARP-Scoping Pitfall:** Simply binding the bypass route to the interface (`route add -host 162.159.192.1
         -interface en0`) fails when the default route is deleted or redirected to `utun*`. Since `162.159.192.1` is
          not local to the physical link, macOS fails to resolve it via ARP, dropping all tunnel packets.
828    - **Fix:** Dynamically resolve the active physical gateway IP (e.g. `172.17.0.1`) on startup, and add the
         static host routes for the WARP endpoints via the gateway IP directly (`route add -host 162.159.192.1 <
         gateway_ip>`). Additionally, because standalone `tailscaled` on macOS does not automatically modify the
         system default route via the CLI, we manually redirect the default gateway to `utun*` using `
         setup_exit_node_routing`.
829
830 2. **The Docker Compose Subnet Takeover Collision:**
831    - **Mechanism:** To avoid collisions with the host's campus network, Colima's default bridge (`docker.bip`)
         is moved (e.g. to `10.200.0.1/24`). However, because `172.17.0.0/16` is now free inside the VM, Docker
         Compose's IPAM dynamically takes it over for the project's custom network (`nullexit_default`). Because the
          containers receive IPs on `172.17.0.0/16`, the host Mac cannot send local peer discovery packets (e.g.,
         Tailscale `magicsock` packets on `172.17.0.2:41641`) directly. The host routing table sends them out `en0`
         to the physical Wi-Fi, returning `host is down` or `no route to host`.
832    - **Fix:** Explicitly configure a custom default network subnet `10.200.1.0/24` in `docker-compose.yml` to
         prevent Docker Compose from ever reclaiming the conflicting `172.17/172.18` ranges.
833
834 > See also 15.2.1 (The Hotspot Paradox) for the related infinite loop that occurs when the physical upstream
         router is also a Tailscale exit-node client.
835
836 #### 15.2.3 Docker Default Bridge vs Host Wi-Fi Subnet Collision (The 172.17.0.0/16 Subnet Overlap)
837 **Context:** When attempting to ping a local Windows client (`172.17.22.187`) or the default gateway
         (`172.17.0.1`) directly over Wi-Fi, the Mac host returns `No route to host` (displaying a `REJECT` / `!`
         flag in the routing table). Tailscale on the client also drops offline shortly after starting when using
         the exit node, failing to establish direct P2P connections and falling back to slow DERP relays.
838
839 **Root Cause & Subnet Overlap Mechanism:**
840 1. **Docker Default Subnet:** By default, the Docker daemon initializes its default bridge network (`docker0`)
         on the **`172.17.0.0/16`** IP range.
841 2. **Subnet Conflict:** The host's physical Wi-Fi network happens to be assigned the exact same
         **`172.17.0.0/16`** range by the local network router.
842 3. **Routing Clash:** Colima launches a Linux VM on the Mac host to run the Docker engine. Because Docker
         inside that VM creates `docker0` and claims the `172.17.0.0/16` subnet, any packet sent to `172.17.x.x`
         from within the containers (such as Tailscale local discovery) is captured by the VM's internal virtual
         bridge interface (which is down/empty) and is blackholed. On the Mac host, macOS dynamically marks routes
         as `REJECT` (`!`) when local ARP resolution fails, causing local pings to immediately abort with `No route
         to host`.
843 4. **Tailscale Relay Saturation**: Because local peer-to-peer discovery is blackholed, Tailscale falls back to
         public DERP relay servers (e.g. `relay "tor"` in Toronto). When exit-node routing is enabled, all client
         web traffic is routed through the relay. Public relays impose strict rate limits and throttle high-volume
         traffic, resulting in packet drops. This saturation drops Tailscale's control packets, causing the
         connection to time out and showing the client as offline.
844
845 **Fix:**
846 The canonical script for this is `scripts/fix-docker-bridge-collision.sh`. **One command applies the entire fix
         :**
847
848 ```bash
849 bash scripts/fix-docker-bridge-collision.sh        # apply
850 bash scripts/fix-docker-bridge-collision.sh --dry  # preview changes without applying
851 ```
852
853 It auto-detects the host's actual Wi-Fi subnet (no more guessing whether the campus DHCP handed out `172.x`,
         `10.x`, or `192.168.x`), picks a non-conflicting `docker.bip` from a small candidate set (defaulting to
         `172.26.0.1/24` if no overlap is matched), backs up `~/.colima/default/colima.yaml` with an ISO-8601
         timestamp, edits the file **in place** (idempotent  replaces an existing `docker.bip` if present, appends
         under an existing `docker:` block, or creates a fresh block), restarts Colima, rebinds Wi-Fi DHCP, verifies
          the new default route is no longer Docker's bridge, and re-runs both `./toggle.sh` and `scripts/diagnose-
         host-leak.sh` to print the post-fix verdict.
854
855 For reference, the underlying manual fix is: edit the Colima configuration file (`~/.colima/default/colima.yaml
         `) on your Mac to define:
856 ```yaml
857 docker:
858   bip: 172.26.0.1/24   # any /24 that does NOT overlap your Wi-Fi subnet
859 ```
860 Then, restart Colima to apply the new subnet inside the VM:
861 ```bash
862 colima restart
863 ```
864 This forces Docker to use `172.26.x.x` for its default bridge, freeing the physical `172.17.x.x` range and
         letting pings and direct Tailscale peer-to-peer connections flow normally.
865
866 #### 15.2.4 July 4, 2026: The "Network Unreachable" Host Leak & Post-Wake Container Destructions
867 **Symptom:**
868 1. Running `docker compose config` with dummy keys corrupted the user's `.env` file by appending invalid
         WireGuard keys. This caused the `warp` container to become unhealthy on boot, which triggered a full
         teardown of the gateway by `toggle.sh`.
```

```
2. After fixing `.env`, `toggle.sh` successfully booted the gateway, but the host's internet started bypassing the tunnel entirely (a massive host leak). The Cloudflare trace returned `warp=off` and exposed the host's physical ISP IP.
3. Running `recover.sh --post-wake` would violently force-recreate all gateway containers and drop the host's internet connection.

**Root Cause:**
- **Bug 1 (The Host Leak):** In `toggle.sh`, the logic to find the Tailscale `utun` interface used `ifconfig | grep -B4 "inet 100."`. Due to the 4 lines of before-context, this regex was accidentally capturing the interface *above* the actual Tailscale interface (e.g. grabbing `utun4` instead of `utun5`). Because `utun4` had no IP address, the subsequent `sudo route add -net 0.0.0.0/1 -interface utun4` silently failed with "Network is unreachable", leaving the host's default route exposed.
- **Bug 2 (The Post-Wake Destructions):** The health check in `recover.sh --post-wake` relied on `docker compose exec -T warp curl -sf https://www.cloudflare.com/cdn-cgi/trace`. However, the newer `qmcgaw/gluetun` Alpine-based Docker image no longer ships with `curl` pre-installed. The health check failed with `executable file not found in $PATH`, tricking `recover.sh` into thinking the tunnel was completely dead and needed to be forcefully recreated via `docker compose up -d --force-recreate`.
- **Bug 3 (Restart Flags):** Previous refactors incorrectly permitted `toggle.sh restart` instead of strictly enforcing `toggle.sh --restart`.

**Resolution:**
- Cleaned the `.env` file of duplicate dummy entries.
- Replaced the brittle `grep -B4` interface parsing in `toggle.sh` with a strict AWK parser: `awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':'`. This mathematically isolates the exact interface possessing the Tailscale `100.x.x.x` IP address, preventing the `0.0.0.0/1` route from silently failing.
- Changed the health check in `recover.sh` to use `wget -qO- --timeout=3` which is natively installed in the Gluetun image. This stopped the unnecessary destructive recreations of the containers during post-wake events.
- Reverted the `restart` alias in `toggle.sh` to strictly require `--restart`.

#### 15.2.5 July 4-5, 2026: Tailscale Data-Plane Loop (The DERP Relay Deadlock)
**Symptom:**
After bypassing the Tailscale control plane (`192.200.0.0/16`) to fix the "offline" mesh status, the Mac appeared online. However, external peer devices (like the user's phone on cellular) could not successfully ping or SSH into the Mac. `tailscale netcheck` reported `Nearest DERP: unknown (no response to latency probes)`.

**Root Cause:**
While bypassing the control plane kept the daemon connected to the coordination servers, Tailscale's actual peer-to-peer data plane relies on STUN (for NAT traversal) and DERP relays. These relays span ~80 dynamically allocated IP addresses across multiple cloud providers. Because we were forcing `0.0.0.0/1` through the `utun*` exit node, the Mac's own outbound connection attempts to DERP relays were being routed back into the tunnel. To prevent an infinite loop, Tailscale's daemon dropped its own packets when it saw them returning on the TUN interface. This effectively left the daemon deaf and blind to peer connections.

**Resolution:**
- Modified `add_warp_bypass_routes` in `scripts/common.sh` to execute a live API fetch (`curl -s https://login.tailscale.com/derpmap/default`) during startup.
- The script parses out all active DERP relay IPv4 addresses (~80 IPs) and adds a physical host bypass route for every single one of them.
- These IPs are written to a temporary file (`/tmp/nullexit-derp-ips.txt`) so they can be cleanly un-routed by `remove_warp_bypass_routes` when the gateway shuts down. Peer connectivity (ping, SSH, SFTP) was fully restored.

#### 15.2.6 July 5, 2026: Hardcoded WARP Endpoints in docker-compose
**Symptom:**
The bash scripts allowed overriding the default Cloudflare WARP IP endpoints via `.env` variables (`WARP_ENDPOINT_1` and `WARP_ENDPOINT_2`) to establish bypass routes. However, `docker-compose.yml` statically hardcoded `162.159.192.1` for the `warp` container's `WIREGUARD_ENDPOINT_IP`.

**Root Cause & Risk:**
If a user set a custom endpoint in `.env`, the script would successfully bypass the custom IP on the host level. However, Gluetun would still attempt to connect to the hardcoded `162.159.192.1`. Since `162.159.192.1` was no longer explicitly bypassed, its packets would be sucked into the `utun*` exit node, causing an infinite WireGuard routing loop that instantly breaks the gateway.

**Resolution:**
Modified `docker-compose.yml` to dynamically read the environment variable via `${WARP_ENDPOINT_1:-162.159.192.1}`, ensuring both the host routing scripts and the container use the exact same endpoint.

#### 15.2.7 Incident Post-Mortem: WARP Watcher Race vs Post-Wake (July 10, 2026)
**Symptom:** After switching Wi-Fi networks, the host IP appeared as the raw ISP IP (`132.x.x.x`) instead of a Cloudflare/WARP IP. The gateway appeared to complete post-wake recovery successfully in the logs, but nuclear shutdown fired immediately after and killed everything.

**Root Cause:** A fatal race condition between two concurrent processes:

1. **Wi-Fi roam**  `watcher.sh` fires `recover.sh --post-wake`
2. **Post-wake** finds the `warp` container missing/dead  calls `docker compose up -d --force-recreate` (~35 seconds)
```

3. **In parallel**, the WARP Watcher (running as a child of `toggle.sh`) polls the warp container every 5s during failures
4. While the container is being recreated it returns `warp=off` for those 35 seconds 7 consecutive failures **WARP SHUTDOWN fires**, calling nuclear `recover.sh`
5. Nuclear `recover.sh` tears down Tailscale, resets DNS to 1.1.1.1, stops all containers **gateway dead, IP exposed**
6. The post-wake `docker compose up` finishes at ~35s mark, container healthy but the gateway is already gone

The evidence in `output.log`:
```
[2026-07-10T06:07:27Z] NET: ...  bash recover.sh --post-wake
[2026-07-10T06:07:47Z] WARP DOWN  cdn-cgi/trace reports warp=off    watcher starts counting
...
add net 0.0.0.0: gateway utun5    post-wake successfully re-routes at 02:08:22
...
[2026-07-10T06:09:04Z] WARP SHUTDOWN  6 consecutive failures         watcher fires nuclear
```

Fix (July 10, 2026): Added a **WARP Watcher inhibit marker** (`/tmp/nullexit-warp-inhibit.marker`):
- `recover.sh --post-wake` writes the marker at startup **before** any container operations
- The WARP Watcher loop checks for the marker at the top of every iteration; if present it resets `consec_off=0` and sleeps 5s (skips the poll entirely)
- `recover.sh` removes the marker at the very end (on both success and failure via `trap cleanup_inhibit EXIT`)
- This guarantees the watcher never counts failures during the intentional container-down window of a post-wake force-recreate

The marker file path is hardcoded to `/tmp/nullexit-warp-inhibit.marker` in both `toggle.sh` and `recover.sh`.

15.2.8 Incident Post-Mortem: Concurrency Collision between Nuclear Teardown and Post-Wake (July 10, 2026)
Symptom: When the background WARP Watcher triggers a nuclear shutdown (due to a real or simulated tunnel failure), the gateway is completely torn down. However, the network configuration changes made during the teardown trigger a `State:/Network/Global/IPv4` network change event. The LaunchAgent-based network watcher (`watcher.sh`) intercepts this event and concurrently spawns `recover.sh --post-wake` to heal/restore the gateway. This causes `docker compose down` (from nuclear teardown) and `docker compose up` (from post-wake) to run in parallel, resulting in container errors like `dependency failed to start: container warp exited (0)` and the gateway becoming wedged.

Root Cause: A lack of coordination/locking between the lifecycle control scripts. While `toggle.sh` used `/tmp/nullexit-toggle.lock` to prevent concurrent `toggle.sh` runs, `recover.sh` (both nuclear and `--post-wake` modes) did not utilize or check any lock files.

Fix (July 10, 2026): Implemented a shared lifecycle lock file (`/tmp/nullexit-toggle.lock`) across both scripts:
- **`recover.sh` (all modes)**: Checks `/tmp/nullexit-toggle.lock` at startup. If the lock is held by another running lifecycle script (`toggle.sh` or `recover.sh`):
  - In `--post-wake` mode, it logs a skip message to `output.log` and exits cleanly (`exit 0`). This safely prevents the auto-recovery agent from recreating containers during a teardown.
  - In nuclear recovery mode, it exits with an error.
- **`recover.sh` Lock Cleanup**: Cleans up the lock file upon completion using `trap cleanup_recover EXIT` (which checks if the lock file belongs to the current PID).
- **`toggle.sh` update**: The lock check in `toggle.sh` was updated to look for both `toggle.sh` and `recover.sh` in the running processes to enforce proper exclusion.

15.2.9 Incident Post-Mortem: Infinite Auto-Recovery Loop after WARP Watcher Nuclear Shutdown (July 10, 2026)
Symptom: When the background WARP Watcher triggered a nuclear shutdown (due to a tunnel outage), it successfully executed `recover.sh` (nuclear). However, immediately after the teardown completed, the system started spawning `recover.sh --post-wake` and rebuilding the containers again. This created an infinite loop of tearing down and auto-restarting the gateway, keeping the host's internet route permanently wedged.

Root Cause: An active-state marker file leak. The LaunchAgent-based network watcher (`watcher.sh`) only fires `recover.sh --post-wake` when `/tmp/nullexit-gateway-active.marker` is present on disk. While `toggle.sh` (STOP path) correctly deleted this marker, the nuclear `recover.sh` script did not. When the WARP Watcher ran `recover.sh` (nuclear), the containers were stopped but the active-state marker was left on disk. The network changes from the teardown then triggered `watcher.sh`, which saw the marker, believed the gateway was still supposed to be active, and triggered `--post-wake` to start it back up.

Fix (July 10, 2026): Modified the start of `recover.sh` to explicitly delete `/tmp/nullexit-gateway-active.marker` if `POST_WAKE` is `false` (nuclear mode). This ensures that any subsequent network configuration changes made during the teardown are correctly ignored by `watcher.sh` because the marker file is gone.

15.3 MTU, Fragmentation & Throughput

15.3.1 Network Bufferbloat & MTU Optimizations (Double-VPN UDP Crash)
Symptom: When running aggressive UDP bandwidth tests (such as macOS `networkQuality`, which uses QUIC/HTTP3), the `nullexit` tunnel completely crashes or exhibits severe lag (6,000ms+ bufferbloat).

Root Cause: The bug is a cascading MTU mismatch. Tailscale generates a UDP packet. Cloudflare WARP (also WireGuard) receives the packet, wraps it in another layer of encryption, pushing the packet size beyond the standard 1500 byte limit. Unlike TCP, UDP packets cannot be resized dynamically via MSS negotiation, resulting in massive IP packet fragmentation. The Apple Virtualization framework (`vz`), which Colima uses

```
     for bridged networking, silently drops heavily fragmented UDP packets under high load. This blackholes the
     entire UDP upload stream, instantly dropping the connection.
960
961  **Fix (Partial):** To prevent MTU fragmentation for standard web traffic, **TCP MSS Clamping** is strictly
     enforced via the macOS `pf` firewall. By adding `scrub out all max-mss 1160` to `scripts/pf.conf`, macOS
     unilaterally shrinks all outgoing TCP headers to fit comfortably inside the double-encrypted tunnel,
     bypassing Path MTU Discovery blackholes. `TS_DEBUG_MTU=1200` was also added to `docker-compose.yml` to
     minimize internal fragmentation.
962  **Unresolved:** The only way to solve the UDP crash bug natively is to artificially cap the maximum tunnel
     bandwidth using SQM (`fq_codel`). Because a static bandwidth cap would artificially throttle high-speed
     networks, `nullexit` leaves the bandwidth uncapped, optimizing perfectly for TCP while accepting UDP
     fragmentation under extreme edge-case load tests.
963  *Why not Dynamic SQM (Autorate)?* Implementing an EMA/PID-controlled autorate daemon (like `sqm-autorate`)
     requires a constant 100ms polling loop to measure the kernel packet queue, which severely degrades laptop
     battery life by preventing deep C-state CPU idling. Furthermore, macOS's native firewall (`dummynet`) only
     supports `fq_codel` and does not support the newer Linux `CAKE` algorithm, which is the only algorithm that
      offers kernel-native, event-driven `autorate-ingress` (zero polling penalty). Therefore, dynamic SQM is
     computationally hostile to battery-powered macOS hosts.
964
965  #### 15.3.2 Double-Tunneling MSS Clamping (Stalling Bug)
966  **Problem:** Traffic double-wrapped through Tailscale (WireGuard) and WARP (WireGuard) suffered 120 bytes of
     MTU overhead (60 + 60). The default MSS clamp of 1180 was still slightly too large, causing large packets
     to fragment or drop, which resulted in mysterious web page stalls on strict-MSS endpoints.
967
968  **Fix:** Lowered the TCP MSS clamp in `post-rules.txt` to `1120` to guarantee double-tunneled payloads fit
     safely within standard internet MTUs.
969
970  #### 15.3.3 Wi-Fi Edge Packet Loss via QUIC (UDP) Fragmentation
971  **Problem:** Applications that heavily utilize HTTP/3 (QUIC) over UDPsuch as Facebook, Instagram, and Google
     servicesexperience massive stuttering and stalled image loading when the client device is far away from the
      router (weak Wi-Fi signal). The issue did not occur when close to the router.
972
973  **Root Cause:** The `nullexit` gateway uses a double-encrypted tunnel (Tailscale + WARP). To prevent packets
     from fragmenting when entering these tunnels, a firewall rule in `post-rules.txt` uses `--set-mss 1120` to
     safely clamp TCP packets. However, because QUIC runs entirely over UDP, it bypasses the TCP MSS handshake
     completely. QUIC blasts maximum-MTU (1280 byte) UDP packets. The inner `tailscale0` interface (MTU 1200)
     forces these large UDP packets to fragment into two pieces. When the client is on the edge of Wi-Fi range (
     where 5-10% packet loss is common), losing *either* UDP fragment destroys the *entire* image frame. This
     exponential amplification of packet loss stalled QUIC streams.
974
975  **Fix & Trade-offs:** Appended an `iptables` rule to `post-rules.txt` that explicitly blocks `UDP --dport 443`
     (QUIC). When modern web apps detect QUIC is blocked, they instantly fall back to HTTP/2 (TCP 443). Because
     TCP is routed through the MSS clamp, the packets are resized *before* they leave, entirely eliminating IP
     fragmentation and restoring high-speed loading over weak Wi-Fi.
976  *Consequences:* This adds a minor latency penalty (TCP 3-way handshake) compared to QUIC's zero-RTT connection.
      It may also force WebRTC video calls (Google Meet/Discord) to fall back to TCP TURN servers, slightly
     inflating real-time video latency. However, in a constrained double-tunnel mesh, the absolute stability
     gained from eliminating fragmentation vastly outweighs these trade-offs.
977
978  #### 15.3.4 Cellular Networks & PMTUD Blackholes (The 1280 Byte Limit)
979  **Experiment:** Conducted a live packet sweep from the PC-1 host to Phone-1 connected via a cellular network +
     Tailscale DERP relay using `ping -D -s <size>`.
980
981  **Result:**
982  - Packets with a payload of `1252` bytes (+ 28 bytes IP/ICMP headers = **1280 bytes total**) succeeded
     perfectly with ~60ms latency.
983  - Packets with a payload of `1253` bytes (**1281 bytes total**) and above resulted in a silent timeout.
984
985  **Analysis:** The cellular network (or the Tailscale DERP relay) enforces a strict MTU of 1280 bytes.
     Critically, it acts as a "PMTUD Blackhole"  it silently drops packets larger than 1280 bytes instead of
     returning an ICMP "Fragmentation Needed" warning. If the exit node tries to send a standard 1500-byte
     internet packet to the phone over this link, it vanishes, causing web pages to stall permanently.
986
987  **Conclusion:** This empirically validates the absolute necessity of the TCP MSS clamping rule in `post-rules.
     txt`. By artificially clamping the MSS to `1120` (or `1180`), we ensure the TCP payload + headers + double-
     WireGuard overhead never exceeds the strict 1280-byte ceiling of the cellular mesh link.
988
989  #### 15.3.5 Low Gateway Throughput (DERP Relay & MTU Fragmentation)  Fixed
990  **Problem:** The host Mac's internet throughput through the gateway dropped significantly (e.g., from 34Mbps
     raw to 2.1Mbps via gateway). This was caused by a combination of two bottlenecks:
991
992  **1. Tailscale DERP Relay Bottleneck:**
993  Because the host Mac and the Docker container operate behind the same public IP, standard hole-punching for
     Tailscale P2P failed, forcing traffic through Tailscale's NYC DERP relay. Normally, `routing-fix.sh`
     explicitly drops container Tailscale UDP packets destined for local subnets (when `TAILSCALE_ALLOW_LAN_P2P=
     false` in `.env` or auto-detected). We nuke these P2P connections on purpose to prevent the severe macOS `
     gvproxy` SNAT endpoint poisoning issue we faced earlier (see 15.7 SNAT Endpoint Poisoning). However, this
     strict policy unintentionally blocked direct communication between the container and the Mac host itself
     via the Docker bridge network.
994
```

Fix: Tailscale's control plane registers the Mac host using its physical IP (e.g., `172.17.52.223`), not the Docker bridge IP. If this physical IP falls within the dropped local subnets (like `172.16.0.0/12`), the P2P handshake is dropped. To solve this natively, `scripts/common.sh` now features a `write_host_ips()` function that actively grabs all physical IPv4 addresses on the Mac and writes them to `.host_ips`. `routing-fix.sh` dynamically reads this file every 30 seconds and injects priority `RETURN` rules for those exact physical IPs. This guarantees direct container-to-host P2P over the local bridge (bypassing DERP) while continuing to block external LAN devices (like phones) to prevent the 15.7 SNAT poisoning bug.

2. MTU Double-Encapsulation Fragmentation:
The host's Tailscale interface (`utun5`) defaulted to an MTU of 1280. The WARP tunnel inside the container also uses an MTU of 1280. When a full 1280-byte Tailscale packet from the host entered the container and was encapsulated by WARP (adding ~60 bytes of overhead), the resulting packet exceeded 1280 bytes, leading to severe fragmentation and performance degradation.

Fix: In `toggle.sh`'s `setup_exit_node_routing` function, the host's Tailscale `utun` interface MTU is explicitly lowered to `1200` (configurable via `HOST_MTU` in `.env`). This ensures that host-originated packets can comfortably fit inside the container's WARP tunnel encapsulation without fragmenting.

> **Design Note Is direct hostcontainer P2P a Docker bug, or are we "exploiting a VM/host relationship"?**
> Neither. It is expected, well-defined behavior, and the RETURN-rule carve-out is a *defensive de-restriction*, not an exploit. The reasoning:
> - **On native Linux**, the container and host share one kernel; the Docker bridge is a first-class local interface and hostcontainer is *trivially* direct. Nobody calls that an exploit. Our RETURN rules simply reproduce that same local reachability on macOS.
> - **The RETURN rules add nothing new; they *stop blocking* something legitimate.** `routing-fix.sh` broadly DROPs containerLAN Tailscale UDP to prevent the real 15.7 `gvproxy` SNAT endpoint-poisoning bug caused by *external* LAN peers (e.g. a phone on another AP). Tailscale's control plane registers the Mac host by its *physical* LAN IP (e.g. `172.17.52.223`), which falls inside those DROPped ranges so the host, the one peer that is literally the same physical machine, became collateral damage. The `.host_ips` RETURN rules are the minimal, precise exception that re-permits exactly that same-machine path and nothing else.
> - **The genuinely VM-specific wart is `gvproxy` SNAT poisoning (15.7), and we route *around* it, not through a hole in it.** Keeping hostcontainer traffic on the Docker bridge (direct) avoids letting it get NATed through the userspace `gvproxy` endpoint rewriter that confuses Tailscale's peer-endpoint learning. That is the opposite of exploiting the VM boundary it is respecting it.
> - **What genuinely *doesn't* work on a VM host is UDP hole-punching to arbitrary *external* peers** (see 15.13). That is a real limitation of userspace VM networking, and it is why remote peers fall back to DERP unless bridged mode is enabled on a trusted LAN. Direct hostcontainer P2P is the one case that works regardless because both endpoints live on the same physical machine, no hole-punching is needed at all. It was verified `direct 172.17.52.223` even in plain user-mode networking (no `--network-address`).
>
> In short: the DROP is the deliberate policy, the RETURN is a surgical exception for same-machine traffic, and the whole thing is standard bridge networking not a Docker bug and not an exploit.

15.4 DNS Interception & Proxy

15.4.1 DNS Proxy: TCP Wire Format Mismatch (socat / dnsmasq)
Problem: When the Tailscale data plane is unavailable (DERP relay only), the script falls back to a local DNS proxy that forwards queries from host UDP:53 to Docker's TCP:5354 (AdGuard). The `socat` and `dnsmasq` approaches both failed.

- **socat** forwards raw bytes it does not prepend the 2-byte length prefix that DNS-over-TCP requires. AdGuard parses the stream incorrectly and returns garbage.
- **dnsmasq** tries UDP first to reach the upstream (`127.0.0.1#5354`), but Colima's SSH tunnel only forwards TCP. With a single upstream server, dnsmasq doesn't fall back to TCP even with `--timeout=1`.

Solution: Replaced both with a **Python DNS proxy** (`scripts/dns-proxy.py`, ~25 lines) that properly handles the DNS-over-TCP wire format: reads a UDP query from the host, prepends the 2-byte length prefix, sends it over TCP to AdGuard, strips the prefix from the response, and sends it back over UDP.

Architectural Note: Why is the DNS proxy in Python while the SOCKS5 proxy was moved to a compiled Go Docker container?
The SOCKS5 proxy handles heavy data streaming (e.g., video, downloads). Python introduces massive CPU/battery overhead for raw socket streaming, which is why we rely on the compiled Go implementation built directly into the `tailscaled` daemon via `TS_SOCKS5_SERVER=:1080`.
Conversely, the DNS proxy only handles intermittent UDP queries (a few bytes per minute) generated solely by the local host machine. The Python overhead for this is effectively 0.001 Watts. We deliberately kept `dns-proxy.py` in Python because it runs natively on the macOS/Linux hostrewriting it in Go would force users to install a Go compiler (`brew install go`) on their host machine for zero noticeable battery gain. (Note: If this architecture is ever adapted to serve as a public DNS resolver for thousands of users or for massive automated web-scraping clusters, `dns-proxy.py` must be rewritten in Go to survive the concurrent packet flood).

15.4.2 Tailscale MagicDNS Bypassing the AdGuard PREROUTING Intercept
Context: The `routing-fix.sh` script applies an iptables NAT rule (`PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT`) to intercept all incoming DNS queries from connected Tailscale peers and force them into the AdGuard container on port 5335.
Problem: When remote clients (like Android or iOS) have "Use Tailscale DNS settings" turned on, they query Tailscale's MagicDNS IP (`100.100.100.100`). This packet enters the exit node, but the `tailscaled` daemon running *inside* the gateway container intercepts the `100.100.100.100` packet internally. `tailscaled` then generates a brand new, local DNS request to the upstream DNS server to resolve the query. Because this new request is generated *locally* by the daemon (not arriving externally via the `tailscale0` interface),

```
     it bypasses the `PREROUTING` chain completely. The request glides straight out the WARP tunnel to
    Cloudflare/Google, entirely bypassing the AdGuard sinkhole.
**Fix:** The only way to fix this blindspot is to force `tailscaled` to use AdGuard as its upstream. In the
    Tailscale Admin Console, the gateway's Tailscale IPv4 address (e.g., `100.x.x.x`) must be added as the sole
     Custom Global Nameserver, and "Override local DNS" must be enabled. Additionally, to block Android devices
     from bypassing this via Private DNS (DNS-over-TLS), outbound Port 853 traffic is explicitly REJECTED in
    both `routing-fix.sh` and `post-rules.txt`.

#### 15.4.3 Seamless Wi-Fi Roaming & The macOS DNS Wipeout Loophole
**Problem:** The Tailscale, WireGuard, and Colima NAT stacks are natively designed for connectionless roaming
    and will seamlessly survive when the macOS host switches Wi-Fi networks (e.g., roaming from home Wi-Fi to a
     cellular hotspot). However, the internet completely breaks upon network transition. This occurs because
    macOS intentionally wipes out all custom DNS settings (`networksetup -setdnsservers`) and reverts to the
    new network's default DHCP nameserver whenever the active Wi-Fi BSSID changes. Since the Mac is locked to
    the exit node, its DNS requests are swallowed by WARP and dumped onto the public internet, which cannot
    route private DHCP IPs. This results in a total DNS failure.

**Solution:** Implementing a silent background "DNS Watcher" daemon (`nullexit-dns-watcher`) in `toggle.sh`.
    When the gateway starts, it spawns a background polling loop that checks the active Wi-Fi DNS every 30
    seconds. If macOS resets the DNS during a network change, the watcher instantly detects the drift and
    forcefully re-injects `100.100.21.8` via `networksetup`. When the gateway stops, the watcher process is
    cleanly killed. This allows true, seamless roaming across physical networks without ever dropping the
    gateway state.

#### 15.4.4 docker compose exec `\r` Injection
**Fix:** All `docker compose exec` output piped through `tr -d '\r'`.

#### 15.4.5 Stderr Invisibly Swallowed
**Symptom:** Failures were silent due to stderr redirected to `output.log`.

**Fix:** Tailscale wait loop now shows output. Critical commands show errors inline.

### 15.5 Roaming, Sleep/Wake & Auto-Recovery

#### 15.5.1 Gateway Breakage on Lid Close / Wi-Fi Roam  Post-Wake + Post-Roam Auto-Recovery
**Context / Symptom.** Closing the lid (sleep/wake) OR losing + reconnecting Wi-Fi (e.g. in an elevator, caf,
    or roaming between APs) leaves the gateway in a partly-dead state. Symptoms range from a ~30s window of
    dead DNS to a permanent outage that only full `recover.sh` or `toggle.sh` fixes. The root cause is that
    nullexit has **no daemon watching macOS sleep/wake or network-state-change events**  the existing DNS
    Watcher inside `toggle.sh` only runs in-process and is suspended along with the rest of the user shell on
    lid close, so the gateway cannot self-heal after either event.

**Diagnosis.** Two distinct failure modes share the same root gap.

* **(A) Lid close  sleep  wake.** `caffeinate -i` (used by `toggle.sh`) blocks *idle* sleep but **does not
    block forced sleep from lid close**. Closing the lid hard-suspends Colima VM + every container + every
    Tailscale-quit-shells-except-tailscaled. On wake the VM clock needs ~5-30s to resync via NTP, during which
    TLS cert validation fails: Cloudflare WARP `162.159.192.1:2408` rejects the handshake  gluetun's
    healthcheck (`interval: 2s, retries: 15`) eventually flips unhealthy  Docker `unless-stopped` restarts the
    `warp` container (~30s gap). `mDNSResponder`'s in-memory cache still holds pre-sleep entries (stale A
    records for tailnet hosts) so the first dozen DNS queries after wake time out. `tailscaled` on the host
    often keeps its stale DERP relay preference and routes packets through a dead relay for another 30-60s.
* **(B) Wi-Fi roam / loss in elevators, captive portals.** WARP is a UDP tunnel to Cloudflare's anycast edge.
    Carrier NATs and captive-portal networks silently invalidate the UDP binding on roam and on hotspot-bound
    re-association. macOS `networksetup` *should* preserve DNS settings on the same Wi-Fi service across a roam
    , but most captive-portal networks DHCP-replace DNS to the captive-portal resolver **during the captive-
    portal dance itself**, before the user has authenticated  so even DNS hijack survives a clean roam and
    quietly breaks on captive-portal handoff.

**Resolution.** A single **launchd LaunchAgent** (`com.nullexit.wake-recovery`) runs a long-running shell
    daemon (`scripts/watcher.sh`) that listens for both event sources and, on each fire, executes `bash recover
    .sh --post-wake`  a *lighter-touch* recovery mode that's the opposite of the existing `--nuclear` semantics
    : it keeps the gateway live while still refreshing every stale subsystem.

**Deep dive: Apple Power Management Primer (for the unfamiliar).**
macOS exposes several relevant surfaces; the right primitive depends on what you actually need to detect. None
    of these are obvious and none of them have a standard splash screen in the docs.

* **`caffeinate <flags>`** creates I/O Kit power assertions via `IOPMAssertionCreateWithName`. It does NOT
    block forced sleep (lid close, low-battery shutdown, sleep timer firing on a per-set Energy Settings policy
    ). Flags:
  * `-i`  PreventIdleSleep (the only one `toggle.sh` uses; protects against the OS going idle while the user is
     active  but a clamshell close still hard-puts it to sleep)
  * `-s`  PreventSystemSleep (only honored on AC power; the user may be on battery)
  * `-u`  DeclareUserActive (resets the idle timer every 30s by default; equivalent to keeping the cursor
    moving)
  * `-d`  PreventDisplaySleep
  * `-m`  PreventDiskIdleSleep
  * There is NO `-l` (lid-block) flag in any modern macOS. Lid close is a kernel-firmware-level event that
    ignores user-space assertions.
```

```
  * To prove any of this on live hardware: `pmset -g assertions` lists every active assertion with type / process / reason / timeout.
* **`pmset`** is the read-side of the same system:
  * `pmset -g log | grep -E 'Sleep|Wake|DarkWake'` shows the recent timeline.
  * `pmset -g history` is the per-day summary.
  * `pmset -g assertions` is the live assertion list (matches `-i -s -d -u` to processes).
* **`launchd` does NOT have a native "on-wake" hook.** Not in any version of macOS as of Sonoma. The canonical recipe is `StartCalendarInterval` with a sub-minute cadence and let launchd queue missed intervals for replay-on-wake, or use a third-party daemon (Sleepwatcher). For our use case a long-running shell daemon listening to the unified log is more responsive than a polling agent.
* **`com.apple.system.powermanagement.*` Darwin notifications** (`willSleep`, `didWake`, `didDim`, `didUndim`) are the C-level signal. From a shell, they are best reached through the unified log with `log stream --predicate`  see the watcher implementation.
* **`scutil n.watch`** is the canonical CLI for live-following network-state changes. Pairs with `n.add State:/Network/Global/IPv4` to get a single tap that fires on **any** IPv4 link/route/DNS/address change across **all** interfaces. This is what `scripts/watcher.sh`'s network-change listener uses.
* **`com.apple.networkChange` Darwin notification** is the C-level equivalent but has no shell-friendly listener; use `scutil` instead.
* **`SCNetworkReachability`** is the Objective-C API used by SOCKS5/HTTP clients to detect route changes per target. Never a fit for shell scripts.

**The Fix (component-by-component).**

1. **`recover.sh --post-wake`** (new flag, **non-destructive** by design). Adding `--post-wake` to the existing nuclear recovery script inverts the semantics. The default mode still tears down Tailscale, resets DNS to empty, stops caffeinate, runs `docker compose down`, power-cycles Wi-Fi. The new mode does:
  * Skip Tailscale disconnect (we WANT it to keep the mesh connection)
  * **Re-hijack** DNS to the gateway Tailscale IP read from `.gateway_ip` (instead of resetting to empty) applies to both `ACTIVE_SERVICE` and `EN0_SERVICE` because the system resolver is scoped to en0
  * **Refresh** the exit-node preference with `tailscale up --reset --ssh=true --accept-dns=false --exit-node=\"$TS_IP\" --exit-node-allow-lan-access=true` (the exact flags `toggle.sh` START uses). This re-asserts DERP mapping without dropping the mesh
  * Inspect `docker inspect --format '{{.State.Health.Status}}' nullexit-warp-1` and **only** `docker compose up -d --force-recreate warp` if gluetun is unhealthy  this single targeted recreate nudges the UDP NAT binding back to Cloudflare without disturbing Tailscale or AdGuard
  * Run the sharingd-reset step (existing fix from `toggle.sh` START path; see 15.11) so AirDrop doesn't freeze on the new IP lease
  * Verify the gateway with `dig +tcp @127.0.0.1 -p 5354 google.com` (AdGuard via TCP through Colima's SSH tunnel) and a 4-second `tailscale ping` to the gateway Tailscale IP, instead of the default mode's direct-to-1.1.1.1 check
  * Crucially: **skip `sudo -v`** because the LaunchAgent has no TTY to prompt on; all privileged calls use `sudo -n` and lean on `/etc/sudoers.d/nullexit` NOPASSWD entries (`networksetup`, `dnsmasq`, `dscacheutil`, `killall`, `pkill`).
2. **`toggle.sh`** writes and clears `/tmp/nullexit-gateway-active.marker` (an ISO-8601 UTC timestamp) at the bottom of the START path and the top of the STOP path / `cleanup_handler`. Two new helpers: `write_gateway_active_marker` and `clear_gateway_active_marker`. The watcher uses this file as its **only signal** that the gateway is live: if `toggle.sh` was stopped (or never started), the watcher no-ops. This prevents waking-up-only-on-counterfactual events from churning DNS.
3. **`scripts/watcher.sh`** (long-running daemon, sibling-style source file). Two backgrounded pipelines:
  * **Wake listener**: `log stream --predicate 'composedMessage CONTAINS[c] "didWake" OR composedMessage CONTAINS[c] "Wake from Sleep" OR composedMessage CONTAINS[c] "Waking from" OR composedMessage CONTAINS[c] "DarkWake"' --style compact --info` piped through `while read; do run_recover "WAKE: ..."; done`. `[c]` makes the predicate case-insensitive (the unified log writes phrases in mixed case across versions).
  * **Network listener**: `(echo "n.add State:/Network/Global/IPv4"; echo "n.watch"; while true; do sleep 86400; done) | scutil 2>/dev/null` piped through a `case` match on `n.state*` / `SCEventUpdate*`.
  * **`run_recover`** checks the marker, debounces via `/tmp/nullexit-watcher.last-recovery` (default 10s, configurable via `NULLEXIT_DEBOUNCE_SECONDS`), and shells out to `bash recover.sh --post-wake`. The 10s debounce prevents a single wake event (which typically emits 2-3 log lines) from triggering multiple recoveries.
  * `trap ... TERM INT`  `kill 0` cleanly tears down the process group on launchctl `bootout` so the `sleep 86400` inside the scutil pipe also dies.
  * `wait` blocks indefinitely while both pipelines feed it.
4. **`launchd/com.nullexit.wake-recovery.plist`** (LaunchAgent in the user domain  runs as the console user at every login without needing sudo).
  * `Label: com.nullexit.wake-recovery`
  * `ProgramArguments: ["/bin/bash", "<absolute-path-to-your-nullexit-install>/scripts/watcher.sh"]`
  * `RunAtLoad: true`  starts immediately on `launchctl load` (no need to log out and back in)
  * `KeepAlive: { SuccessfulExit: false, Crashed: false }`  **don't** restart on clean SIGTERM exit (so `launchctl bootout` doesn't get stuck in a relaunch loop). Also **don't** restart on hard crash: we observed that macOS Sonoma+'s sleep/wake suspend-resume path accumulates orphan watcher.sh PIDs because SIGCONT does not reliably clean up the prior instance's listener grandchildren, and a `KeepAlive.Crashed=true`-driven relaunch actively races with the still-live prior process. The script enforces single-instance on its own via a `flock`-style PID-file lock, so a relaunch-on-crash would be skipped by the lock and produce 0 listener coverage until the live instance happened to die. Trade-off: a real crash leaves the user without auto-recovery until they run `launchctl load -w` manually  acceptable because (a) gateway still works without auto-recovery, (b) a relaunch wouldn't fix whatever caused the crash, and (c) the gateway-recovery itself is observably broken (no DNS, no exit-node) if it does go down.
  * `ProcessType: Background`  hint to App Nap not to suspend us.
  * `StandardOutPath/StandardErrorPath: /tmp/nullexit-watcher.{out,err}.log`  separates launchd-level diagnostics from the script's own `/tmp/nullexit-watcher.log` (which `exec >>` redirects).

```

Install (one-time). The plist lives in the repo at **`launchd/com.nullexit.wake-recovery.plist`** for version control. The installed (live) copy must land in `~/Library/LaunchAgents/` so launchd picks it up on the user session.

```bash
# Copy the plist to the user-launchd location (run from repo root)
cp ./launchd/com.nullexit.wake-recovery.plist \
   ~/Library/LaunchAgents/

# Substitute __NULLEXIT_HOME__ with your local install absolute path.
# The plist uses WorkingDirectory + a relative program arg, so this
# single substitution is the only place it references your machine.
sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" \
    ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Validate the XML
plutil -lint ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Load + persist this user session + every future login
launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist

# Confirm it actually started
launchctl list | grep nullexit
```

To **stop** the watcher (without uninstalling):

```bash
launchctl bootout gui/$UID/com.nullexit.wake-recovery
# or:
launchctl unload -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```

To **uninstall** completely:

```bash
launchctl bootout gui/$UID/com.nullexit.wake-recovery 2>/dev/null || true
rm ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
```

Why this fix was preferred over alternatives.

* **Polling with `StartCalendarInterval`** would have given ~30-60s detection latency per wake. The unified-log stream is sub-second.
* **Sleepwatcher** would have covered wake events but not network changes. We'd still need a separate `scutil` listener, and introducing a third-party dependency for what amounts to ~150 lines of shell is unjustified.
* **Forcing `caffeinate -l`** doesn't exist and the closest equivalent (`caffeinate -s`) only works on AC power. The whole point of the gateway is to stay usable on battery while traveling lid close must NOT silently kill the gateway.
* **A kernel extension** is impossible (no entitlement) and out of scope.

Reproduction recipe. Test both events in isolation to confirm the fix actually closes the gap.

* **(A) Lid-only repro**: with the gateway up, run in a terminal: `pmset displaysleepnow && sleep 30 && pmset wake` (without physically closing the lid). Watch the wake wake-fires the watcher post-wake routine runs login should still work in <5s after wake. Compare to baseline pre-fix: pre-fix the gateway would be dead for 30-90s.
* **(B) Roam-only repro**: with the gateway up: `networksetup -setairportpower off && sleep 30 && networksetup -setairportpower on` (or simply walk out of Wi-Fi range and back). The captive-portal WILL repave DHCP DNS for a few seconds; the watcher fires on the second `State:/Network/Global/IPv4` change (after the Wi-Fi re-associates) and the post-wake routine re-hijacks DNS within 2-3s.

Operative files.

* `recover.sh` added arg parsing; wrapped destructive sections behind `POST_WAKE`; added re-hijack DNS / refresh exit-node / force-recreate-if-unhealthy path.
* `toggle.sh` added `write_gateway_active_marker` / `clear_gateway_active_marker` called at the END of START and TOP of STOP / cleanup_handler.
* `scripts/watcher.sh` new long-running daemon.
* `launchd/com.nullexit.wake-recovery.plist` launchd LaunchAgent descriptor.
* `~/Library/LaunchAgents/com.nullexit.wake-recovery.plist` the installed copy (one-time copy).

Honest assessment (read before testing). This fix has two asymmetric halves:

* **Network / roam / captive-portal path RELIABLE.** `scutil n.watch` on `State:/Network/Global/IPv{4,6}` catches every Wi-Fi roam, captive-portal DHCP-rebind, hotspot handoff, and VPN change with zero missed events. To validate, drop Wi-Fi for 30 s and re-add it; `/tmp/nullexit-watcher.log` will record a `NET:` trigger and `recover.sh --post-wake` will run within ~10 s of the network coming back.
* **Lid close / wake path BEST-EFFORT on macOS Sonoma+.** Apple *suspends* LaunchAgents process-state during forced sleep and *resumes* them on wake; it does NOT re-launch them every wake. As a result: (a) `log stream` is a live subscription events emitted during the agent's suppressed window are DROPPED, not

```
     buffered, so the wake event that prompted the resume is invisible to `log stream` on resume; (b) the "fire
     on startup" guard runs once on `launchctl load -w` and after a `KeepAlive` crash-recovery, NOT on every
     successive wake; (c) `subsystem == "com.apple.powermanagement"` will catch FUTURE emissions (display dim/
     restore, manual sleep/wake) but not the lid-closeopen wake directly.
1163
1164 In practice the lid-wake gap is short, because `toggle.sh`'s existing DNS Watcher already polls every 5 s and
     re-hijacks DNS (so DNS recovers within 5 s of any roam or wake), `tailscaled` self-heals via DERP fallback
     within 3060 s, and `gluetun` reconnects via its healthcheck retry within 30 s. The user-visible pain on lid
      open is ~30 s of wonky DNS, not a permanent outage.
1165
1166 **Test the ROAM path FIRST.** If ROAM works in your testing, the lid-wake gap is acceptable. If lid-wake
     handling is critical, the two clean follow-ups are:
1167 1. **Sleepwatcher** (one canonical install): `brew install sleepwatcher`; drop `recover.sh --post-wake` into
     `~/.sleep` and `~/.wake` so Sleepwatcher invokes it directly on every wake without the launchd propagation
     problem.
1168 2. **Move the watcher to a LaunchDaemon** (root) and use `IOPMSchedulePowerEvent` via a tiny C wrapper  not
     portable to shell.
1169
1170 A truly reliable shell-only solution to "wake events from inside a LaunchAgent" does not exist on macOS Sonoma
     +.
1171
1172 **Empirical lessons from deployment.** Two findings came up while deploying this fix end-to-end; recording so
     the next operator doesn't lose an evening to each.
1173
1174 1. **Bash function-call-before-definition gotcha.** While introducing the gateway-active marker helpers, the
     defensive `clear_gateway_active_marker` call was inserted at the very top of `toggle.sh` (right after `set
     -e`), but the function definition lived near `PID_FILE="/tmp/nullexit-caffeinate.pid"` ~90 lines down. Bash
      parses top-to-bottom and resolves names at first invocation, so the call site exited with code 127 (`
     clear_gateway_active_marker: command not found`). The gateway stayed unreachable from the START path even
     though `bash -n` passed (the parser doesn't catch order-of-resolution errors). Rule: define functions
     BEFORE any block that calls them. Verify with `grep -n '^name()\s*{'` followed by `grep -n '^name\s*$'`
     the latter must appear on a line number AFTER the former.
1175 2. **`networksetup -setairportpower off+on` is not a faithful roam reproduction.** Synthetic Wi-Fi radio-
     cycling (`sudo networksetup -setairportpower Wi-Fi off` for 30 s, then back on) is expected to produce a `
     NET:` trigger in `scutil n.watch` but produced ZERO during testing  because `setairportpower` powers the
     radio at the driver level while leaving the network SERVICE in the `"up"` state from scutil's perspective,
     so no `n.state State:/Network/Global/IPv4` event is generated. A real roam (macOS briefly loses the AP and
     re-associates into a different one) DOES fire the event because the link actually changes.
1176    Better reproduction recipes in priority order:
1177    * Walk between two real SSIDs via `networksetup -switchtolocation` (or actual physical station movement)
       guaranteed to fire the scutil event.
1178    * Force a DHCP rebind on the same SSID with `sudo ipconfig set en0 DHCP`  triggers `State:/Network/Global/
       IPv4` to update.
1179    * Trust the existing 30-second DNS Watcher inside `toggle.sh` for the DNS path: it re-hijacks DNS
       automatically. The post-wake watcher adds clear value mostly for tailscale exit-node re-assertion and warp
       gluetun force-recreate, both of which self-heal within ~30 s anyway.
1180
1181 #### 15.5.2 July 8, 2026: Network Watcher Event Mismatch, Sudo-in-Background Crash, and Loop-Prevention
1182 **Symptom:**
1183 1. Switching Wi-Fi or waking the Mac from sleep did not trigger a post-wake recovery (`recover.sh --post-wake`)
       automatically.
1184 2. The network connection would eventually get completely sinkholed/blocked. If the user tried to restart or
     wait, it did not self-heal.
1185 3. During the recovery attempt, checking the public IP on the host would return the unencrypted campus/ISP IP (
     starting with `132.`) instead of the encrypted tunnel IP.
1186
1187 **Root Cause:**
1188 - **Bug 1 (Network Watcher Mismatch):** In `scripts/watcher.sh`, the network change listener watched `State:/
     Network/Global/IPv4` changes via `scutil n.watch` but matched them against `*n.state*|*SCEventUpdate*`. On
     macOS, the actual output is `changedKey [0] = State:/Network/Global/IPv4`. Because the pattern did not
     match, all network switches were silently ignored.
1189 - **Bug 2 (Sudo Caching Background Crash):** When the background WARP Watcher eventually detected the broken
     tunnel (after 6 failures/30s), it triggered the default nuclear `recover.sh` (with `POST_WAKE=false`).
     Because there was no active TTY in the background context, the `sudo -v` caching check aborted instantly
     with a terminal required error. Due to `set -e`, `recover.sh` died immediately before resetting DNS,
     disabling the packet-filter killswitch, or stopping containers, leaving the host network completely
     sinkholed.
1190 - **Bug 3 (Post-Wake Infinite Loop):** Once the watcher patterns were fixed, `recover.sh --post-wake` triggered
      successfully on network changes. However, because the script deleted the default routing entries at the
     start, spent ~15 seconds rebuilding 88 DERP relay bypass routes, and re-added the default routes at the end
     , this execution time exceeded the watcher's 10-second debounce threshold. Furthermore, the watcher
     recorded the *start* timestamp in the debounce file. Consequently, the route changes made by the script
     itself triggered the watcher again immediately after completion, trapping the system in an infinite loop of
      recovery runs. During this loop, the VPN routes were constantly deleted, resulting in the host leaking
     traffic through the real ISP interface (an IP starting with `132.`) for 15s out of every cycle.
1191
1192 **Resolution:**
1193 - **Fixed Watcher Matching:** Updated `scripts/watcher.sh` to match `*changedKey*` and `*State:/Network/Global/
     IPv*` to correctly catch network changes.
1194 - **Bypassed Background Sudo:** Added `--auto`/`--non-interactive` flags to `recover.sh` and added a TTY check
      `[ -t 0 ]` to skip interactive `sudo -v` authentication when running headlessly in the background. Updated
```

```
    the WARP Watcher call in `toggle.sh` to pass `--auto`.
- **Prevented Infinite Loops:** Modified `scripts/watcher.sh` to write the **completion** timestamp of `recover.sh` to `DEBOUNCE_FILE` (instead of only the start timestamp). This ensures all buffered network config events generated during route reconstruction are evaluated against the end-of-run timestamp and successfully debounced, breaking the infinite loop.
- **Centralized & Timestamped Logs:** Redirected `watcher.sh` stdout/stderr from `/tmp/nullexit-watcher.log` to the repository's main `output.log` and updated `scripts/common.sh`'s `step`, `ok`, `warn`, `fail`, and `die` helpers to prefix logs with a local `[YYYY-MM-DD HH:MM:SS]` timestamp.

#### 15.5.3 Incident Post-Mortem: Startup Pre-Flight Check Tailscale Race (July 10, 2026)
**Symptom:** Immediately after startup via `toggle.sh`, the pre-flight connectivity check `[1/3] Gateway reachable via Tailscale` returned `FAIL`, forcing `toggle.sh` to skip the full exit-node activation and fall back to SOCKS5/local DNS proxy mode.

**Root Cause:** A timing race condition during startup. In `toggle.sh`, the host joins the mesh using `tailscale up --reset` (Phase A) right before performing the pre-flight check. Because `tailscale up` resets and reconfigures the host's `tailscaled` daemon, the local daemon must fetch the latest netmap asynchronously from the coordination servers. This network map propagation and route establishment takes a few seconds. Running `tailscale ping` immediately after `tailscale up` returns (with no delay) causes the check to fail because the local daemon has not yet discovered the gateway node `100.100.21.8`.

**Fix (July 10, 2026):** Upgraded pre-flight Check 1 in both `toggle.sh` and `scripts/toggle-linux.sh` to use a 5-attempt retry loop with a 1-second delay between checks. This allows up to 5 seconds for local tailnet routing paths to settle, ensuring the gateway is correctly reached and a pong is received before deciding whether to enable exit-node routing.

#### 15.5.4 Incident Post-Mortem: Pre-Flight Check False Negatives due to VPN Firewall STUN Blocks (July 10, 2026)
**Symptom:** Immediately after a cold boot or full restart of the gateway, the pre-flight connectivity check `[1/3] Gateway reachable via Tailscale` failed, causing `toggle.sh` to skip the full exit-node configuration and fall back to SOCKS5/local DNS proxy mode. However, manually running `tailscale ping 100.100.21.8` immediately after startup succeeded in 1ms.

**Root Cause:** An asynchronous Tailscale handshake delay caused by firewall restrictions. Inside the container network namespace, the `warp` container (gluetun) implements a strict firewall that blocks all outbound UDP traffic to the public internet except to the designated WireGuard endpoint. Because of this, the `tailscale` container's STUN requests to public STUN servers are blocked, causing the local daemon to report `UDP is blocked, trying HTTPS`. Without UDP STUN capability, Tailscale is forced to fall back to a slower DERP-assisted handshake (relayed via HTTPS/TCP). This negotiation and direct-path punch-through takes roughly 30 to 35 seconds to establish on cold boots (after a full `tailscale up --reset`). Since the pre-flight check only waited 15 seconds (15 attempts with 1-second sleeps), the check timed out and aborted before the path completed.

**Fix (July 10, 2026):** Increased the pre-flight ping retry limit from 15 to 45 attempts (45 seconds) in both `toggle.sh` and `scripts/toggle-linux.sh`. This provides sufficient time for the DERP-assisted handshake to complete on cold boots, while still passing immediately (in 1-2 seconds) on warm starts or once the path is established.

#### 15.5.5 Incident Post-Mortem: Docker Image Build DNS Failures on Immediate Restart (July 10, 2026)
**Symptom:** When executing `toggle.sh --restart`, the script successfully stops the gateway but then immediately fails during startup during `docker compose build` with DNS resolution errors:
`failed to do request: Head "https://registry-1.docker.io/...": dial tcp: lookup registry-1.docker.io on 192.168.5.1:53: no such host`.

**Root Cause:** A race condition between physical link negotiation and virtual machine startup. The STOP path of the restart command invokes `cleanup_network_state` which power-cycles the host's physical Wi-Fi interface (`en0`) to flush stale routing states. Immediately after, `toggle.sh` starts the gateway boot sequence. Because `toggle.sh` did not verify the status of the physical interface, it proceeded to boot Colima and build Docker containers while `en0` was still negotiating a DHCP lease. Consequently, the host (and by extension, the VM bridge) had no active internet gateway/DNS path, causing Docker's registry check to fail.

**Fix (July 10, 2026):**
1. Created a unified `wait_for_dhcp_settle` helper function in `scripts/common.sh` that polls `ipconfig getpacket` and `ifconfig` until a valid IP and router are assigned. To handle typical macOS Wi-Fi reassociation and DHCP negotiation latency (which commonly takes 15-20 seconds after a power-cycle), this loop was configured to poll up to 60 times (30 seconds total).
2. Injected a call to `wait_for_dhcp_settle` at the beginning of the `START` action in `toggle.sh` (right before checking Colima status) to block container builds until the host's physical network is online.
3. Refactored `recover.sh` to reuse the unified `wait_for_dhcp_settle` function.

#### 15.5.6 Watcher.sh Suppressed Exit Code Bug (Fixed)
**Observation:** The `scripts/watcher.sh` script is a long-running daemon that invokes `recover.sh --post-wake` when it detects system wake or network roam events. Previously, it contained the following bug:
```bash
  bash "$RECOVER" --post-wake
  local exit_code=$?
  ...
  return $exit_code || true
```
```

```
1231 The `|| true` mistakenly swallowed the actual `$exit_code` returned by `recover.sh`, causing `run_recover()` to
       always return `0` (success), masking any underlying failures in the wake recovery process.
1232
1233 **Fix:** The `|| true` statement was removed from the return line:
1234 ```bash
1235   return $exit_code
1236 ```
1237 Since `watcher.sh` explicitly omits `set -e` (to prevent pipe breakages from killing the daemon), removing `||
       true` safely allows the underlying exit code to propagate without risking daemon termination. The watcher
       daemon was then reloaded (`launchctl kickstart -k`) to pick up the script changes in memory.
1238
1239 #### 15.5.7 Network Readiness Race Condition on `--restart` (Fixed)
1240 **Problem:** When running `./toggle.sh --restart`, the script tears down the gateway and immediately begins the
       startup sequence. On systems with Wi-Fi (or any wireless interface), the OS takes a moment to re-associate
       with the network after the teardown's DNS/routing changes. If the startup sequence ran before the OS had a
       default route, `get_active_service` would return nothing, routing bypass targets would be wrong, and DNS/
       WARP setup would silently fail  resulting in a partially configured gateway that self-heals only once the `
       watcher.sh` re-runs.
1241
1242 **Root cause:** The original code had no gate before the first network-dependent call (`ACTIVE_SERVICE=$(
       get_active_service)` at the top of the start branch). This was a pure timing race.
1243
1244 **First fix (Wi-Fi specific):** A polling loop on `ipconfig getifaddr en0` was added to wait for an IPv4
       address on `en0` before proceeding. This worked but was brittle  it only handled Wi-Fi and would fail for
       Ethernet, USB-C adapters, or iPhone USB tethering.
1245
1246 **Generalized fix:** Replaced the `en0` check with `route get default | grep 'gateway:'`. This detects a valid
       default gateway on *any* interface, covering all connection types. The loop polls every 2 seconds with a
       60-second hard timeout (then proceeds with a warning rather than hanging forever). The output prints the
       detected gateway IP so failures are easy to diagnose in logs.
1247
1248 ```bash
1249 # In toggle.sh, before ACTIVE_SERVICE=$(get_active_service)
1250 while true; do
1251   if route get default 2>/dev/null | grep -q 'gateway:'; then
1252     _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
1253     echo "  Network ready (default gateway: $_gw)."
1254     break
1255   fi
1256   ...
1257 done
1258 ```
1259
1260 **Why not ping?** Pinging an IP (e.g. `1.1.1.1`) during restart is unreliable because DNS may still be hijacked
       from the previous session, and the exit node routing may not yet be cleared, causing the ping to loop back
       through the tunnel. `route get default` is a pure local kernel query  no packets sent.
1261
1262 ### 15.6 WARP Tunnel Liveness & Host Leaks
1263
1264 #### 15.6.1 Host-Side Traffic Leaking Past the Exit Node (with Remote Clients Routing Correctly)
1265 **Symptom.** `toggle.sh` reports success. Remote tailnet clients (e.g. an S24 phone) correctly egress through
       Cloudflare WARP and see a Cloudflare IP on `whatismyip.com`. But on the **HOST Mac running the gateway**, `
       whatismyip.com` reports the underlying physical ISP's ASN (e.g. `campus_isp` for a university campus Wi-Fi)
       . The gateway container itself is healthy; the leak is specifically on the host.
1266
1267 **Root causes.** Three distinct mechanisms can each produce this exact symptom on the host alone, while leaving
       remote clients unaffected. They are mutually rankable so the diagnostic picks the active one
       deterministically.
1268
1269 * **(A) SOCKS5 fallback lane active.** `toggle.sh`'s three Phase-B pre-flight checks (`tailscale ping`  AdGuard
       via `dig +tcp`  WARP internet, toggle.sh ~lines 750-820) sometimes flake during the first minute. When ANY
       of the three fails, the script sets `SKIP_EXIT_NODE=true` (toggle.sh:849) and instead activates the SOCKS5
       fallback path (toggle.sh:894). The SOCKS5 fallback hijacks DNS to `127.0.0.1` (so AdGuard filtering still
       works) but **modern browsers on macOS do NOT honor the system SOCKS5 proxy**  they ignore `networksetup -
       setsocksfirewallproxy`. Result: DNS gets ad-filtered but actual HTTP/S traffic exits direct through the
       campus ISP. This branch is invisible from a remote tailnet client because the client's tunnel is unaffected
       .
1270
1271 * **(B) IPv6 leak over campus Wi-Fi.** The Tailscale exit node and WARP tunnel are both IPv4-only (gluetun + `
       post-rules.txt` drop all IPv6 forwarded traffic; see 15.12 IPv6 Exit-Node Leak). But many university and
       enterprise APs are dual-stack  the host happily accepts AAAA DNS responses and sends IPv6 traffic straight
       out `en0`. macOS happily encrypts that traffic through Tailscale ONLY if Tailscale has IPv6 advertised;
       with an IPv4-only exit node, IPv6 traffic skips Tailscale entirely. This branch is invisible from a phone
       because iOS/Android Tailscale apps actively sinkhole IPv6 in their VPN profile, but the macOS host's `
       tailscaled` does not.
1272
1273 * **(C) Route-table freeze and `--accept-routes` reset after `tailscaled` wake/toggle.** Running `tailscale up
       --reset` clears the exit-node preference and reverts `--accept-routes` back to its default (`false`).
       Without explicitly passing `--accept-routes=true`, `tailscaled` will not attempt to install the default
       route through the `utun*` tunnel interface even if the exit node preference is reconciled. Additionally,
       macOS occasionally freezes and fails to apply route-table changes (see 15.11 Standalone Tailscale Daemon
```

```
     Freeze). In both cases, `netstat -rn` shows the default route still pointing to the physical Wi-Fi gateway (e.g. `172.17.0.1`) instead of the Tailscale `utun*` interface, causing host traffic to exit direct while remote clients route correctly.
1274
1275 **Why remote clients don't see this.** Each remote phone/laptop uses its OWN Tailscale tunnel to reach the container's `100.x.x.x:443` over the mesh  they're end-to-end encrypted regardless of how the host Mac itself is configured. Only the HOST's own egress sockets (HTTP requests from the host's browser, curl, etc .) are affected.
1276
1277 **Resolution.** A single script: `scripts/diagnose-host-leak.sh`. It runs the 8 checks, classifies into one of A / B / C / OK, prints a verdict + a ready-to-run fix command on stdout, AND writes the full report to `host-leak-diagnostic-<UTC>.txt` so the user can paste the file back instead of scrolling-and-copying from the terminal. With `--fix`, the matched remediation is applied and the two checks that could change (warp + ipv6) are re-verified. With `--watch` (or `--watch 30`), it runs the full baseline diagnostic then enters a continuous loop checking warp/IPv6/default-route every N seconds, alerting only on state changes  logs to `host-leak-watch.log`.
1278
1279 ```bash
1280 # Diagnose only:
1281 bash scripts/diagnose-host-leak.sh
1282
1283 # Diagnose + apply matched fix + re-verify:
1284 bash scripts/diagnose-host-leak.sh --fix
1285 ```
1286
1287 **What it prints.** A 7-section report (`tailscale status`, SOCKS5 state, `cdn-cgi/trace` warp=, IPv6 leak probe, host default route, last toggle.sh output.log hints, resolved gateway Tailscale IP), followed by a classification block (`Scenario: A | B | C | OK`) with the exact command(s) the user needs to paste into their terminal  they don't have to figure out which `tailscale up` flags match their environment.
1288
1289 **What it does NOT do.** It does not touch `toggle.sh` or `toggle.sh`'s pre-flight logic, does not modify `.env` (except `--fix` mode appends `GATEWAY_BYPASS_PING=true` when fixing Scenario A), and does not restart Docker. It is a read-only diagnostic with an opt-in remediation flag. Safe to run on a healthy gateway  the worst case is a 30-line report that says "OK".
1290
1291 **Why cdn-cgi/trace over whatismyip.com for the verdict.** `whatismyip.com`'s ISP database routinely mislabels campus IPv6 ranges as residential ISPs (that's why the local ISP's ASN appears even when you're routing through Cloudflare). `cdn-cgi/trace` is hosted by Cloudflare ITSELF and reports `warp=on|off` plus the exact edge colo serving the request. It is the only public endpoint that can truthfully confirm whether the WARP tunnel is in use.
1292
1293 **Common false alarms.**
1294
1295 * **`warp=on` but whatismyip.com still shows the local ISP.** `whatismyip.com`'s ISP-detection database is unreliable on academic networks. Trust `cdn-cgi/trace`'s `warp=on` over `whatismyip.com`'s ISP string. Use `https://api.ipify.org` (returns just the IP, no ISP DB lookup) as a third confirmatory check.
1296 * **`tailscale status` shows the host online but exit node preference is missing.** `tailscale up --reset --exit-node=` (empty value) clears any stale preference; `tailscale up --reset --exit-node=<100.x.x.x>` re-establishes it. The diagnostic prints the exact command with the resolved TS_IP filled in.
1297 * **Both IPv6 leak AND route-freeze flags set simultaneously.** Scenario B outranks Scenario C  fixing IPv6 makes IPv4 the only viable egress, after which the route-freeze usually self-resolves within ~30s (or one more `toggle.sh` cycle).
1298
1299 **Deep dive: Host Leak Probe  Continuous Sub-Second Egress Monitor (`scripts/host-leak-probe.sh`).**
1300
1301 *What it is.* `scripts/host-leak-probe.sh` is a lightweight background daemon (~1.4 MB RSS, ~01% CPU) that polls `https://www.cloudflare.com/cdn-cgi/trace` every 300ms directly from the host via `curl`  not via `docker compose exec`. This is a fundamentally different vantage point from the WARP Watcher (15.6.2): the WARP Watcher asks the WARP *container* whether it sees `warp=on`; the Host Leak Probe asks what a real browser request leaving the host's physical NIC looks like. The two blind spots it covers that the WARP Watcher cannot:
1302
1303 1. **Host-side flash-leaks**  a Cloudflare anycast edge re-route, a browser reusing a stale pooled connection across a tunnel state change, or a split-second routing gap during a Gluetun healthcheck restart (see 15.12.13 Routing Fix 30-second polling  acknowledged 14s window).
1304 2. **Host egress while container is healthy**  the container can report `warp=on` while the host's own traffic exits direct (the three-scenario leak class documented above in 15.6.1).
1305
1306 *Lifecycle.* Started by `toggle.sh`'s START path (`start_host_leak_probe`) immediately after `start_warp_watcher`. Controlled by:
1307
1308 | Signal path | What happens |
1309 |---|---|
1310 | `toggle.sh` STOP | `stop_host_leak_probe()`  SIGTERM + PID file cleanup |
1311 | `toggle.sh` `cleanup_handler` (error/INT/TERM/HUP) | Same |
1312 | `recover.sh` (default, non `--post-wake`) | 6d block  kill by PID file |
1313 | `recover.sh --post-wake` | Left running  the gateway is still live |
1314
1315 PID file: `/tmp/nullexit-host-leak-probe.pid`. Toggle with `HOST_LEAK_PROBE=false` in `.env` to disable at next gateway start.
1316
```

```
*Output format.* All events are written to `output.log` (repo root) with UTC timestamps. Only state *changes* are logged  the probe is silent during normal healthy operation.

```
[09:10:26] LEAK warp=off ip=45.23.1.1 prev_warp=on prev_ip=104.28.246.50
[09:10:26] ROTATE warp=on ip=104.28.248.11 prev_ip=104.28.246.50
[09:10:26] HOST-PROBE failed/timeout (count=1)
```

curl transport errors (`TLS handshake failed`, `Connection refused`, etc.) are also appended to `output.log` via `2>>output.log` nothing is silenced to `/dev/null`.

Grep cheatsheet.

```bash
grep LEAK output.log            # real warp=off events from the host
grep ROTATE output.log          # Cloudflare anycast edge rotations (harmless)
grep 'HOST-PROBE' output.log    # curl failures / timeouts
```

Resource budget. ~1.4 MB RSS, ~0% CPU at rest, ~3 HTTPS requests/second to Cloudflare. Safe to leave running for hours. Back off the polling interval (`bash scripts/host-leak-probe.sh 1.0`) if you observe probe-failure pile-ups indicating Cloudflare rate-limiting.

Detection vs prevention. This script detects leaks; it does not prevent them. A sub-300ms flash-leak can still slip between two polls undetected. If `grep LEAK output.log` shows confirmed leak events, the next step is a kill-switch (prevention) an `iptables` rule on the host's `en0` that drops non-Tailscale, non-local egress whenever the gateway is active. That is a separate engineering task not yet implemented.

15.6.2 In-Flight WARP Tunnel Liveness Monitor & Statistical Auto-Shutdown

Why. nullexit's core promise is that your IP always appears as Cloudflare. If the WARP tunnel silently dies due to a UDP NAT timeout, a Cloudflare edge rotation, a Colima VM clock skew causing TLS cert rejection, or any of a dozen other failure modes the host silently falls back to egressing through the physical ISP. The user sees the gateway as "up" (containers running, Tailscale connected, DNS hijacked) but their real IP is exposed. This is the exact class of failure that Ingo Blechschmidt documented in his chilling 2024 Linux kernel post-mortem: for over two years, LUKS disk encryption on suspend was silently broken because a refactored kernel syscall returned success while doing nothing ([source](https://mathstodon.xyz/@iblech/116769502749142438)). The lesson: **a security mechanism that fails silently is worse than no mechanism at all** it creates a false sense of safety that prevents the user from taking corrective action.

Design principle. The WARP Watcher (`start_warp_watcher()` in `toggle.sh`) is a background daemon that polls `cdn-cgi/trace` every 30 seconds and applies a **statistically calibrated threshold** before triggering any safety response. This threshold distinguishes transient network blips (which self-heal within seconds) from genuine outages (which require intervention). A single `warp=off` reading is meaningless Docker might be slow, the network might be congested, a container healthcheck might be restarting. But 6 consecutive off readings (30 continuous seconds of downtime) has vanishingly low probability of being a false positive: even at an unrealistically high 10% false-positive rate per poll, P(6 consecutive) = 0.1 0.0001%. In practice, the per-poll false-positive rate is far lower than 10%, making the real probability astronomically small.

What it does.

1. **Always logs.** Every state transition (onoff, offon) is timestamped to `output.log`. The user can `grep WARP output.log` to audit the tunnel's entire lifetime. This is the "never fail silently" mandate.

2. **Tracks consecutive failures.** A counter increments on each `off` reading and resets to 0 on any `on` reading. This ensures the watcher never fires on intermittent flapping.

3. **Auto-shuts down at threshold.** When the counter reaches `WARP_FAIL_THRESHOLD` (default 6, configurable in `.env`), the watcher declares a statistically significant outage and runs `recover.sh` the nuclear recovery script that disconnects Tailscale, stops Docker containers, resets DNS to `1.1.1.1`, flushes routes, and power-cycles Wi-Fi. This instantly reverts the host to normal (un-encrypted) internet, guaranteeing no traffic leaks through a dead tunnel while the user thinks they're protected. A trigger file at `/tmp/nullexit-warp-shutdown-triggered` prevents double-firing.

4. **Exits cleanly.** After shutdown, the watcher exits (the gateway is dead, there's nothing left to monitor). `stop_warp_watcher()` is also called by `toggle.sh`'s STOP path and `cleanup_handler` to terminate the daemon cleanly during normal teardown.

Configuration. Set `WARP_FAIL_THRESHOLD=3` in `.env` for a more aggressive 15s timeout, or `WARP_FAIL_THRESHOLD=12` for a conservative 60s window. The variable is parsed from `.env` using the same `grep` pattern as `GATEWAY_MSS` and `GATEWAY_BYPASS_PING`.

Log sample. After a real outage (or a forced test via `docker compose pause warp`):

```
[2026-07-03T21:15:17Z] WARP DOWN  cdn-cgi/trace reports warp=off
[2026-07-03T21:15:47Z] WARP SHUTDOWN  6 consecutive failures (threshold=6). Running recover.sh to kill the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.
[2026-07-03T21:15:52Z] WARP SHUTDOWN  recover.sh completed, watcher exiting. Your IP is now your real ISP IP.
```
```

```

Or, if WARP self-recovers before the threshold:

```
[2026-07-03T21:15:17Z] WARP DOWN cdn-cgi/trace reports warp=off
[2026-07-03T21:15:32Z] WARP RECOVERED warp=on after 3 consecutive off readings (threshold was 6)
```

**Relationship to other monitors.** The WARP Watcher is the **innermost safety layer**  it runs as a child of `toggle.sh` and only while the gateway is active. It complements (does not replace) the `scripts/watcher.sh` daemon (15.5.1, which handles sleep/wake and Wi-Fi roam) and `scripts/diagnose-host-leak.sh` (15.6.1, which is a manual diagnostic tool). The WARP Watcher is the only component that actively verifies end-to-end tunnel liveness and takes autonomous action when it fails.

#### 15.6.3 Incident: The Overnight Silent IP Leak (July 2026)
**The Problem:** The host leaked traffic over its raw, unencrypted `eth0` IP continuously for roughly 8 hours, completely bypassing the VPN. The `toggle.sh` state and the container liveness checks all reported the gateway as healthy and active.

**The Investigation:**
1. **Sleep/Wake Checks:** We initially suspected macOS sleep cycles broke the routing table. A check of `pmset -g log` confirmed the laptop literally never slept (0 sleep events recorded since boot), ruling out all sleep/wake code paths.
2. **Keepalive Expiry:** We discovered `docker-compose.yml` was missing `WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL`. Without a keepalive, standard router NAT tables (which typically garbage collect UDP after 30 to 120 seconds of silence) quietly expired the WireGuard port mapping overnight.
3. **Container State Masking:** Because WireGuard is stateless, the `warp` container didn't immediately crash. It stayed "Up", which meant Docker's healthchecks and our `toggle.sh` liveness checks never noticed the tunnel inside it was totally dead.

**The Solution:**
1. Added `WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL=25s` (the WireGuard standard to beat standard 30s NAT timeouts) to keep the UDP tunnel permanently pinned open through the router.
2. Injected a **Fail-Closed Kill-Switch** into `routing-fix.sh`: A 30-second loop now actively verifies tunnel health by running `curl` over `tun0` to Cloudflare's trace endpoint. If it fails 3 consecutive times, it immediately rewrites the default route to `blackhole`, severing all traffic instead of letting it leak through `eth0`.
3. Added a listener in `watcher.sh` that detects this fail-closed event and instantly pushes a native macOS UI notification to the desktop, alerting the user to manually intervene.

> [!NOTE]
> **Why dual failure systems?**
> The system now relies on two independent failure handlers that cover each other's blind spots:
> 1. **Automated Self-Healing (WARP Watcher + `recover.sh`):** Triggered when the `warp` container logs `warp=off`. This happens when the server actively kicks the client. Because the container *knows* it is disconnected, it triggers `recover.sh` to safely reboot Docker and reconnect.
> 2. **Fail-Closed Kill-Switch (`routing-fix.sh` + `watcher.sh`):** Triggered when a silent network failure occurs (e.g., NAT timeouts). The container incorrectly believes it is connected (`warp=on`), so auto-recovery never fires. Instead of auto-recovering (which risks falling back to raw Wi-Fi mid-process while the user isn't looking), this kill-switch actively blackholes the internet and forces a manual intervention via a desktop notification.

### 15.7 Tailscale P2P, SNAT & Control Plane

#### 15.7.1 macOS SNAT Endpoint Poisoning & The P2P Blackhole (July 10, 2026)
**Symptom:** When a Tailscale client (like an S24) connected to the exact same Wi-Fi network as the macOS gateway, it completely lost internet connectivity. However, if the S24 connected to a Windows PC's hotspot on the same network, the tunnel worked flawlessly.

**Root Cause:** Claude's critique correctly dismantled the port collision theory: the gateway daemon wasn't failing to bind, and macOS wasn't load-balancing ports. The real culprit was **macOS SNAT Source Port Mangling poisoning Tailscale's path discovery**.

Here is the exact step-by-step of the "infinite loop" blackhole:
1. **The P2P Handshake:** The S24 and the Mac are on the same Wi-Fi (`192.168.1.x`). The S24 discovers the Mac's local IP and sends a UDP hole-punch packet to `192.168.1.5:41641`. Colima successfully forwards this to the gateway container.
2. **The Bypass Rule:** The gateway replies. Because `routing-fix.sh` forces all Tailscale UDP traffic out `eth0` to bypass WARP, the reply goes to the macOS host.
3. **macOS SNAT Mangling:** macOS routes the reply out `en0` to the S24. Because the packet crosses from the Colima VM bridge to the Wi-Fi interface, macOS applies Source NAT (SNAT). Critically, macOS **randomizes the source port** (e.g., changes it from `41641` to `58291`).
4. **Endpoint Poisoning:** The S24 receives the authenticated Tailscale packet from `192.168.1.5:58291`. Tailscale's `magicsock` is designed to handle NAT dynamically, so it assumes the Mac has roamed to port `58291`! The S24 updates its endpoint and starts sending all its encrypted internet traffic to `192.168.1.5:58291`.
5. **The Blackhole:** macOS has no port forward for `58291`. It drops 100% of the S24's outbound data packets. The connection is completely dead.
```

```
**Why did the Hotspot work?** When the S24 connected to the PC hotspot, it was behind the PC's NAT. The S24 sent the packet, PC NATted it, and the Mac replied (mangled to `58291`). When the reply hit the PC, the PC's strict NAT dropped it because the source port (`58291`) didn't match the destination port it originally sent to (`41641`). Because the mangled packet was dropped, the S24 never poisoned its endpoint! It safely gave up on P2P, fell back to the 100% reliable TCP DERP relay, and the internet worked perfectly.

**Fix & Resolution (July 10, 2026):**

**Phase 1  Port disambiguation:** Changed the gateway container's Tailscale listen port from the default `41641` to `41642` across `docker-compose.yml`, `post-rules.txt`, and `routing-fix.sh`. This prevents any bind conflict with the macOS host's native Tailscale daemon.

**Phase 2  RFC1918 DROP rules:** Updated `routing-fix.sh`'s `add_tailscale_p2p_bypass()` to explicitly drop all outgoing Tailscale UDP packets destined for private subnets (`192.168.0.0/16`, `10.0.0.0/8`, `172.16.0.0/12`) when `TAILSCALE_ALLOW_LAN_P2P=false`. This surgically kills the local P2P hole-punch attempt *before* it can trigger the macOS gvproxy SNAT mangling. The S24 receives a clean failure and immediately locks onto the public DERP relay with 0% packet loss.

**Phase 3  Automatic network detection:** Because `TAILSCALE_ALLOW_LAN_P2P=false` breaks P2P on trusted hotspots that genuinely don't have AP isolation, a `.env` toggle and a full auto-detection system were implemented:

- **`.env` toggle:** `TAILSCALE_ALLOW_LAN_P2P=true/false`  static fallback, used only if auto-detection cannot run.
- **`watcher.sh` auto-detection (`detect_lan_p2p_mode`):** On every network change, `watcher.sh` runs two checks:
  1. **WPA2-Enterprise detection:** `airport -I` is used to read the current Wi-Fi `link auth` field. If the auth type does not contain `-psk` (e.g., it is plain `wpa2` = 802.1x), the network is classified as enterprise and P2P is forced off. *(Note: the `airport` binary was removed in macOS Sonoma 14.4  the detection now uses `system_profiler SPAirPortDataType`; see 15.11 for the full migration.)*
  2. **AP Isolation probe:** For WPA2-Personal networks, a short `ping` is sent to the default gateway. If the gateway responds (it always should on a real home/hotspot network), P2P is allowed. If not, AP isolation is inferred and P2P is forced off.
  - The result (`true` or `false`) is written to `.lan_p2p_detected` in the repo root.
- **`routing-fix.sh` dynamic reading:** `add_tailscale_p2p_bypass()` re-reads `.lan_p2p_detected` on every 30-second loop iteration and atomically adds or removes the RFC1918 DROP rules. **No container restart is needed when switching networks.**

**Known Unknowns (Pending Verification):**
- The gvproxy SNAT source port mangling theory is mechanistically sound and backed by observed behavior (gvproxy's UDP proxy changelog has documented edge cases in this exact code path), but has not yet been confirmed via `tcpdump -i en0 udp portrange 40000-60000` while triggering a P2P handshake from the phone. A packet capture cross-referenced with `tailscale ping` endpoint reports on the phone side would confirm the theory definitively.
- Claude's recommendation: run `tcpdump -i en0 udp port 41642 or portrange 40000-60000` on the Mac while triggering the handshake from the phone to see whether the reply leaves with a different source port than `41642`.
- **Client-Side VPN Cycling (Roaming Recovery):** An observation (July 10, 2026) suggests that when a hung DNS probe state is triggered by roaming between Access Points (APs) on a WPA2-Enterprise network, simply toggling the VPN connection (Tailscale) off and on locally on the client phone immediately resolves the issue. You do not need to cycle the phone's physical Wi-Fi. This reinforces the theory that roaming on Enterprise networks leaves a stale/poisoned P2P endpoint in the phone's Tailscale magicsock state, which gets instantly flushed upon a local VPN restart. (Status: Possible but not formally verified).

#### 15.7.2 July 6, 2026: Packet Filter (pfctl) Defaulting Local LAN Rules to TCP-Only
**Symptom:** Devices on the exact same local Wi-Fi network as the gateway were experiencing ~500ms latency to the gateway instead of the expected ~80ms.

**Root Cause:** When writing the `pf.conf` rules to whitelist local LAN traffic (`pass out quick on $ext_if to { 192.168.0.0/16, ... }`), the rule omitted an explicit protocol. The macOS `pfctl` compiler implicitly applies state tracking to all `pass` rules. However, because state tracking (`keep state`) natively evaluates TCP sessions, `pfctl` automatically appended the `flags S/SA` modifier to the compiled rule in the kernel. This silently restricted the whitelist to **TCP only**, causing all local UDP traffic to be dropped by the default-deny kill-switch. Since Tailscale relies on UDP port 41641 for direct P2P connections, the local devices were unable to hole-punch and were forced to fall back to routing traffic through a remote Tailscale DERP relay, massively inflating latency.

**Resolution:** Split the single implicit rule into explicit `proto tcp`, `proto udp`, and `proto icmp` rules in `scripts/pf.conf` to guarantee the macOS compiler allows all three core transport protocols on the local subnet without mutating the rule into a TCP-only lock.

#### 15.7.3 Incident Post-Mortem: PF Kill-Switch Bricks Host Internet in SOCKS5 Fallback Mode (July 10, 2026)
**Symptom:** When the pre-flight checks fail (e.g. because host Tailscale was unauthenticated/logged out), the script falls back to SOCKS5 proxy mode but the host immediately loses all internet connectivity and tailscaled reports `You are logged out... connect: no route to host`. Even running `sudo tailscale up` to log in fails because the connection to the control plane times out.

**Root Cause:** An architectural mismatch. The macOS Packet Filter (PF) Kill-Switch works at the IP level: it blocks all outbound traffic on the physical interface (`en0`) except to explicitly whitelisted VPN endpoints, expecting all other traffic to go through the virtual VPN interface (`utun*`). In SOCKS5 fallback mode, the exit-node (`utun*`) is NOT enabled; instead, only specific applications (like browsers)
```

use the localhost SOCKS5 proxy, while all other host traffic (including the `tailscaled` daemon's UDP packets) continues to go through `en0`. Because the script still enabled the PF Kill-Switch during SOCKS5 fallback, it blocked all non-SOCKS5 traffic on `en0`, completely bricking the host's internet and locking the Tailscale daemon out of the control plane.

Fix (July 10, 2026): Modified `toggle.sh` to remove the `enable_killswitch` call from the SOCKS5 fallback path. The firewall kill-switch is now strictly reserved for exit-node VPN mode, ensuring SOCKS5 fallback leaves the host's direct internet route open for system daemons to authenticate.

15.7.4 Incident Post-Mortem: Host Tailscale Logouts due to Aggressive reset Flags (July 10, 2026)
Symptom: Every time the user runs `toggle.sh` or `recover.sh` which disconnects/reconnects the host client, the host client randomly gets logged out, forcing them to re-run `sudo tailscale up` to authenticate.

Root Cause: An overly aggressive configuration reset. The host-side `tailscale up` calls (used in `disconnect_tailscale_host` and the startup/enable exit node paths) all passed the `--reset` flag. On macOS, `--reset` resets the entire configuration state and authentication token cache of the daemon. Since the script does not pass a `TS_AUTHKEY` to the host's commands (which are meant for the user's private interactive client), the `--reset` flag effectively logged the user out of their tailnet, requiring interactive re-login.

Fix (July 10, 2026): Removed the `--reset` flag from all host-side `tailscale up` invocations in `toggle.sh` and `scripts/common.sh`. This ensures that exit-node state transitions (enabling/disabling) are handled safely while preserving the host client's authenticated session.

15.7.5 Incident Post-Mortem: macOS Tailscale Flag Requirement Abort (July 10, 2026)
Symptom: When the `toggle.sh` stop trap or the roaming watcher triggered, the gateway failed to shut down or disconnect properly. The `output.log` showed the error: `Error: changing settings via 'tailscale up' requires mentioning all non-default flags. To proceed, either re-run your command with --reset...`

Root Cause: In a previous attempt to stop host-side logouts (15.7.4), we removed the `--reset` flag from all `tailscale up` invocations. However, if the Tailscale daemon on macOS is already running with non-default flags (like `--ssh` or `--accept-routes`), invoking `tailscale up` with only a subset of flags (like `--exit-node=`) triggers a strict safety check in the Tailscale CLI that aborts the command completely. Because the command aborted, exit-node routing remained stuck.

Fix (July 10, 2026): Migrated all specific preference changes in `toggle.sh` and `scripts/common.sh` from `tailscale up` to a two-step sequence:
1. `tailscale up` (with no flags) to safely ensure the interface is brought up using the current authenticated state.
2. `tailscale set --exit-node=...` to modify the specific target preference cleanly without triggering the missing flags error or requiring `--reset`.

15.8 Docker / Colima Lifecycle

15.8.1 Colima VM Memory / Swap Thrashing Latency
Symptom: After running nullexit flawlessly for over 24 hours straight to test stability, the 0.5GB VM RAM configuration began experiencing severe network latency on the Tailscale mesh later in the day.

Root Cause: Services (like AdGuard, Tailscale, Docker logs) slowly accumulate cache and state over extended periods. Under the tight 512MB physical RAM limit, this slow creep forced the Linux kernel to aggressively swap memory to the 512MB SSD swap file. The constant disk I/O thrashing blocked process execution, causing DNS resolutions and packet forwarding to delay significantly (making the internet feel "very slow").

Proof (Empirical Observation): While in this OOM-thrashing state, direct DNS queries through the Tailscale mesh (`dig @100.100.21.8 google.com`) would intermittently hang or take several seconds to resolve. After restarting with the 600MB configuration, the exact same query immediately dropped to a lightning-fast **41 ms** response time. (Note: Querying the mapped host port via `127.0.0.1:5354` will always time out due to Gluetun's strict leak-prevention firewall blocking non-VPN ingress traffic).

Fix: Increased VM base physical RAM to 600MB (`--memory 0.6`) and reduced the swap file to 400MB. This provides enough native RAM headroom for long-running state caches without forcing heavy I/O swapping. Subdomain deduplication still reduces blocklists by ~60%. Memory profiles (`light`/`medium`/`heavy`) let users trade off coverage for memory.

15.8.2 Missing iptables in routing-fix Container
Root Cause: `alpine:3.20` does not have iptables installed. Commands failed silently with `2>/dev/null || true`.

Fix: Updated `docker-compose.yml` to `apk add --no-cache iptables iproute2` before `scripts/routing-fix.sh`.

15.8.3 DOCKER_GW Variable Parsing Bug
Root Cause: `ip route show default` printed two lines; `awk '{print $3}'` picked up both, causing route commands to fail.

Fix: Added `head -1` to the parsing chain.

15.8.4 Hard Reboots Leave Stale Docker Socket in Colima VM
Problem: If the macOS host machine is subjected to a hard reboot, sudden power loss, or a kernel panic, the Colima VM running in the background is abruptly killed. Upon next boot, running `toggle.sh` resulted in the contradictory output:

`Colima is already running.` followed by `Cannot connect to the Docker daemon at unix://~/.colima/default/docker.sock.`

Root Cause: The hard crash prevents the inner Docker daemon from executing its graceful shutdown routines. It leaves behind a corrupted `docker.sock` file inside the VM. On boot, the outer VM spins up (hence "already running"), but the inner Docker daemon sees the stale socket, assumes another instance is running, and crashes.

Fix: Added an Auto-Healing routine directly into `toggle.sh` (Step 4b). The script now explicitly tests the Docker daemon with `docker ps`. If the daemon is unresponsive despite Colima reporting as active, it assumes a stale socket and automatically runs `colima restart` in the background to cleanly rebuild the VM and socket before proceeding.

15.8.5 Cloudflare WARP 24-Second Startup Delay (UDP Session Rate-Limiting)
Context: When running `toggle.sh` to turn the gateway ON, `docker compose up -d` often hangs for exactly 24-25 seconds waiting for the `warp` container to become `healthy`. Users might mistakenly blame the Go `rule-compiler` processing 400k+ rules for this delay, as Docker prints `rule-compiler Exited 24.3s` at the end of the wait.
Root Cause: The `rule-compiler` actually finishes its work in <2 seconds. The 24-second blockage is entirely due to Cloudflare WARP's edge server rate-limiting. WireGuard is a stateless protocol with no "disconnect" packet. When the toggle is turned OFF, the old UDP session state remains active on Cloudflare's backend for 20-30 seconds. When the toggle is instantly turned back ON, Docker starts a new WireGuard connection from a new ephemeral UDP source port but uses the **same static cryptographic key** (from `.env`). Cloudflare detects this as a potential replay attack or key-sharing abuse, and intentionally drops all packets (rate-limits) for the new connection until the old session's ghost state fully expires (~20-25 seconds).
Result: Gluetun's healthcheck fails repeatedly every 2 seconds until Cloudflare lifts the penalty box, at which point traffic flows and Docker unblocks the rest of the startup sequence. This is a known, expected behavior of reusing static keys rapidly on Cloudflare's network and cannot be bypassed.

15.8.6 July 4, 2026: Multi-stage Docker Builds for Go Ports & Teardown Hang Fix
Symptom: The Python implementations of `logger.py` and `sync-rules.py` were slow and required a heavy python runtime inside Alpine containers. Also, `routing-fix` container teardown was hanging for 30s during `toggle.sh` shutdown.
Root Cause:
- Python is slower than Go and installing it dynamically via `apk add` added overhead and complexity to the containers.
- Docker containers ignore `SIGTERM` if the init process doesn't handle it, causing a full 30s timeout on shutdown.
Resolution:
- Ported `logger` and `sync-rules` to Go using Goroutines for massive concurrent speedups.
- Implemented **Multi-stage Docker builds** (using `golang:1.22-alpine` as builder) to compile the static binaries inside Docker. The final containers are raw Alpine with zero dependencies.
- Added `--build` to `docker compose up -d` in `toggle.sh` to ensure updates are consistently applied on boot.
- Added `stop_grace_period: 1s` to `routing-fix` in `docker-compose.yml` which eliminates the 30-second teardown hang instantly.
- Implemented a cryptographic integrity checker (`scripts/crypto.sh`) using HMAC-SHA256 and `NULLEXIT_SEED` in `.env` to prevent tampering of bash scripts. (See also 15.9.4.)

15.8.7 Colima `start` Post-Provision Hang (~2 min) & the Docker-Readiness Poll (July 14, 2026)
Symptom: `toggle.sh` took ~197s of wall-clock time to bring the gateway up, and almost the entire delay lived in a single step: `colima start`. Timestamps in `output.log` showed the VM reaching `READY` and creating the `colima` Docker context in ~8-12 seconds, after which `colima start` simply *did not return* for ~110 more seconds until the `run_with_timeout 120` watchdog `SIGKILL`ed it. The rest of the boot (containers, routing) then proceeded normally.

Root Cause: A colima-internal hang that occurs *after* `Successfully created context colima` is printed. It is **mode-independent**. It was initially misattributed to `--network-address` / `--network-mode bridged` needing the `socket_vmnet` helper, but a controlled test with plain user-mode networking (`colima start --memory 0.6 --vm-type vz`, no `--network-address`) reproduced the identical ~120s hang (04:40:34 04:42:34, watchdog-killed) while Docker was fully usable the entire time. Crucially, killing the hung `colima start` process does **not** stop the VM and does **not** corrupt state `colima status` correctly reports `running` afterwards and `docker` keeps working via the `colima` context. The ~110s was therefore pure dead time waiting on a CLI that had already done its real work.

Fix: `scripts/common.sh` now provides `colima_start_until_ready()`. It launches `colima start "$@"` in the background and polls `docker --context colima info` every 3s. The moment Docker answers (~12s in practice) it stops waiting, reaps the (usually still-hung) starter with `kill`, runs `docker context use colima` to guarantee the default context is set, and returns 0. If Docker never answers within a 90s cap it returns non-zero and `toggle.sh` continues to the existing Step 4b `docker ps` auto-heal (15.8.4). Both Colima start call sites in `toggle.sh` the LAN-P2P bridged/shared path and the user-mode fast path now go through this helper instead of `run_with_timeout 120 colima start`.

Result: The Colima phase dropped from ~120s to ~12s and total `toggle.sh --restart` wall time from ~197s to ~106s, with the double-tunnel and direct hostcontainer P2P confirmed intact after the change. The key insight: treat "Docker is reachable" not "the `colima start` CLI returned" as the definition of *started*.

15.8.8 Toggle-On Startup Optimizations: Rule-Compiler Off the Critical Path (July 14, 2026)
Context: A genuine cold toggle-on (VM bounced for RAM, WARP down) measured **95s**. Profiled per-phase with the now-fixed `DEBUG_TRACE` (15.11.8-A): Colima cold start **19s** (near floor), `docker compose up` **45s

, tailscale connect+poll **21s, final checks **7s**.

Two wrong hypotheses first (documented so they aren't re-tried): (1) that BuildKit `--build` was the ~30s cost it wasn't; every layer was already `CACHED`, so gating `--build` saved only ~2s. (2) that `docker compose exec` was slow in the poll loop it isn't (0.18s steady-state vs 0.06s for `docker exec`); the ~5s/iteration was `tailscale status` **blocking during cold connect** and hitting the `run_with_timeout 5` cap.

Real bottleneck (evidence-backed): the 45s `compose up` was **not** WARP. On a genuine cold start WARP goes healthy fast (tor starts right with it no 15.8.5 same-key penalty; that only bites rapid offon). The long pole is the **rule-compiler re-deduping ~340k rules inside the 600MB VM (~28s), every toggle** `adguardhome` started 28s after warp because it waited on `rule-compiler: service_completed_successfully`, even though all 16 remote lists loaded from cache unchanged.

Fix hash-gated, off the critical path (kept in-container):
- `rule-compiler` is now a **`compile`-profiled** service, excluded from `docker compose up`. `adguardhome` and `routing-fix` no longer `depends_on` it they boot instantly on the `compiled_rules.txt` / `ip_blocklist.ipset` already in `./adguard/work`.
- `toggle.sh` runs it on-demand via `docker compose run --build --rm rule-compiler`, **gated on `compute_rules_hash()`** = sha256 of `black_list.txt` + `white_list.txt` (the lists the user controls). Recompile only when that hash changes, when an artifact is missing, or when the compiled file is older than `RULES_MAX_AGE_DAYS` (default 7, so the remote blocklists still refresh occasionally). A compile failure is **non-fatal** we keep the previous compiled rules rather than bricking the tunnel (ad-block the kill-switch).
- Also trimmed the tailscale-connect poll timeouts (status `5s3s`, `ip -4` `10s5s`) to reduce overshoot.

Why NOT run the compiler natively on the Mac host (the tempting idea): rejected on purpose. There is no Go on the host, so "native" means either `brew install go` (a host toolchain dependency that drifts the project from *portable+secure-on-any-device* toward *faster-on-mine*) or committing a prebuilt arch-specific **binary blob** you'd then have to trust/sign. Worse, it would make the **host** write into the `./adguard/work` bind-mount violating the deliberate *"single allowed writer is the rule-compiler container"* invariant (hostcontainer UID/perms mismatches corrupt AdGuard's BoltDB store). Staying in-container preserves both properties; the gate delivers the speed.

Result: unchanged toggle `compose up` **45s 9s**, total `--restart` **~95106s 63s**, with AdGuard still serving the full 342,519 rules (it reuses the existing compiled file) and the double-tunnel + direct hostcontainer P2P intact. Editing `black_list.txt`/`white_list.txt` flips the hash one-off recompile new rules take effect. Both lists were added to the `crypto.sh` signed set (they're security inputs), so editing them requires a re-sign which is the same moment you'd want the recompile. Local state files `.build_hash` / `.rules_hash` are gitignored.

15.9 Permissions, Crypto & Security Hardening

15.9.1 Sudo Credential Caching Removed
Problem: The script used `sudo -v` at startup to cache sudo credentials, plus a background `SUDO_KEEPER_PID` loop refreshing them every 60 seconds. This required interactive password entry even with NOPASSWD sudoers configured, because `sudo -v` itself prompts.

Fix: Removed the entire `sudo -v` + `SUDO_KEEPER_PID` section. All privileged commands now use `sudo -n` (non-interactive), which works silently because the user has NOPASSWD rules in `/etc/sudoers.d/nullexit` for the specific commands needed (`networksetup`, `python3`, `dscacheutil`, `killall`, `pkill`, `route`, `ifconfig`).

15.9.2 Background Sudo Failures
Context: The `--post-wake` script is intentionally designed to skip the `sudo -v` interactive prompt because it is launched by `watcher.sh` running via `launchd` in the background (which has no TTY and cannot accept password input). Thus, it relies entirely on `sudo -n` (non-interactive sudo) for all its internal routing changes.
Problem: If the user has not configured the `/etc/sudoers.d/nullexit` file properly, any automated background wake event will silently fail: `sudo -n` commands drop out, the WARP bypass routes fail to apply, and the `warp` container loses internet connectivity.

15.9.3 Tailscale Hijacking the Default Route (PHYSICAL_GW Bug)
Problem: When `--post-wake` runs, it clears the routing table using `sudo route -n flush` to ensure a clean slate after the Wi-Fi card roams or wakes up. Because it skips bouncing the Wi-Fi interface (to make the reconnect faster), macOS's DHCP client does not automatically rebuild the physical default route. To counteract this, the script must manually restore the physical default route. Originally, the script tried to capture the physical gateway via `route get default`. However, because the gateway is already running and Tailscale is active, the default route is often already pointing to `utunX`. When this happens, `route get default` does not output a `gateway:` field (it only outputs the `interface:`), causing the captured `PHYSICAL_GW` variable to be empty.
When `PHYSICAL_GW` is empty, the physical default route is never restored, causing the `en0` interface to be isolated from the physical network. The WARP bypass routes then fall back to targeting `-interface en0` (instead of routing via the gateway), which breaks WireGuard's remote connections completely.
Fix: The script now bypasses the OS routing table entirely and directly queries the hardware DHCP lease (`ipconfig getpacket en0`) to reliably extract the physical router IP, ensuring the physical route is correctly restored even when Tailscale is active.

15.9.4 July 4, 2026: Consolidation of Core Bash Scripts (Refactoring)
Symptom: Over 200 lines of bash logic were directly duplicated across `toggle.sh` and `recover.sh` (e.g. `restart_tailscaled_daemon`, `stop_sleep_prevention`, WARP bypass routes). This caused drift when one script was updated without the other.

Root Cause: Historical separation of responsibilities allowed `recover.sh` to grow complex alongside `toggle.sh`.
Resolution:
- Massively refactored both scripts by extracting all duplicated networking logic, PID lifecycle management, and environmental configuration down to `scripts/common.sh`.
- Specifically, the `stop_sleep_prevention`, `stop_dns_watcher`, `stop_host_leak_probe`, and `stop_warp_watcher` functions were collapsed into a single `stop_pidfile_daemon()` abstraction.
- Proxy disable loops were abstracted to `disable_all_proxies()`.
- The `162.159.192.1/.193.1` WARP edge IPs are now dynamically resolved from `.env` instead of hardcoded.

15.9.5 Threat Model: Local Proxy Discovery
The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). The real security boundary here is "can an unprivileged local process reach the loopback interface," not "does the malware know the port number."

The port-randomization and the Honey-Port trap only catch blind or generic port-scanners that do not already know where to look. They **do not protect** against a process that directly reads the `/tmp/nullexit-ports.env` file, greps the `toggle.sh` memory space, or runs `lsof -i` to inspect open socket bindings.

Because modifying file ownership on `/tmp/nullexit-ports.env` via `chown`/`chmod` to restrict read-access introduces an unavoidable Time-Of-Check to Time-Of-Use (TOCTOU) race condition (the file is created user-owned before privilege escalation, meaning file descriptors can be snatched), it is deliberately left as a standard user-owned file.

The actual, proper mitigation for preventing a malicious local process from reaching the proxy is not hiding the portit is running a host-level outbound firewall with strict localhost/loopback filtering enabled (e.g., LuLu 4.3.0+ or Little Snitch). These firewalls gate access by cryptographic process identity (binary signature) at the kernel/network-extension level, rather than by port secrecy.

Verifying Localhost Filtering (The Rogue Binary Test). To empirically validate that the host firewall properly intercepts local proxy hijacking, you can compile and execute a completely unknown, un-whitelisted "rogue" binary to attempt a connection to the proxy port:

```c
// lulu-test.c
#include <stdio.h>
#include <stdlib.h>
#include <arpa/inet.h>

int main(int argc, char *argv[]) {
    int port = atoi(argv[1]);
    int sock = socket(AF_INET, SOCK_STREAM, 0);
    struct sockaddr_in server = { .sin_family = AF_INET, .sin_port = htons(port) };
    server.sin_addr.s_addr = inet_addr("127.0.0.1");

    printf("Rogue binary attempting connection to 127.0.0.1:%d...\n", port);
    if (connect(sock, (struct sockaddr *)&server, sizeof(server)) < 0) {
        printf("Connection BLOCKED by firewall.\n");
        return 1;
    }
    printf("Connection SUCCESSFUL (Firewall bypassed!).\n");
    return 0;
}
```
Compile with `gcc -o lulu-test lulu-test.c`. When executed against the `$TOR_SOCKS_PORT`, a properly configured host firewall (with "Allow Loopback" disabled in LuLu Preferences) will immediately pause the socket execution and generate a desktop alert for the unrecognized binary signature. This physically stops targeted local malware from exfiltrating data through the Tor tunnel, proving the threat model mitigation.

15.9.6 Unlocking Mode-Locked Output Files Without `chmod` (`unlock-files.sh`)
The one-command fix: `bash scripts/unlock-files.sh` replaces every known locked inode in the project with a fresh writable one. No `chmod`. Covers `.env` (mode `000`), `adguard/work/userfilters/compiled_rules.txt`, `ip_blocklist.ipset`, `cache/*.txt`, `cache/ip/*.txt`, and `adguard/work/data/filters/*.txt` (mode `0444`). Run it once after setup or after pulling an old repo; `toggle.sh` and `sync-rules.py` will then work without permission errors.

Context: The nullexit project has a strict project-wide policy of `no chmod from scripts` (see README 6). This rule exists because every prior `chmod 0444` (post-write tamper-proof lock) and `chmod 000` (post-toggle `.env` lock) call from `scripts/sync-rules.py` / `toggle.sh` could leave a file permanently unreadable on disk if the script exited early, was killed mid-flight, or simply set a restrictive mode without ever being re-flipped. We've stripped every `chmod` from the codebase. But files **written by older versions of those scripts** are still on disk in those restrictive modes (`0444` for `compiled_rules.txt`, `ip_blocklist.ipset`, `data/filters/<id>.txt`; `000` for `.env` after end-of-toggle lock). Re-running `toggle.sh` or `scripts/sync-rules.py` against those leftovers hits `PermissionError: [Errno 13] Permission denied` immediately.

Why the stale permissions also block `mv`/`rm` of the parent folder: The restrictive modes (`000` on `.env`, `0444` on output files) don't just block writes they prevent the kernel from unlinking the file's dentry during a `mv` or `rm` of the **containing directory** (`adguard/work/`). macOS enforces this at the VFS layer: you own the directory, but you can't delete a directory that contains entries you can't write to. This made the entire `adguard/work/` tree unmovable/undeletable until the locked inodes inside it were

```
     replaced. `unlock-files.sh` resolves this permanently by swapping each locked inode for a fresh `0644` one
     via atomic rename (`cp` + `mv`), which operates on the directory entry rather than the file contents  no `
     chmod` called, no permission bypass needed.
1594
1595 **Problem:** You (the owner) cannot `cat`, `cp`, `rm`, `> redirect`, or any normal POSIX write against a `mode
     =000` file even though you own it  the basic mode bits block ALL access for owner, group, and other. The
     script does not call `chmod` and you want a permanent fix that never relies on a chmod from any future run
     either.
1596
1597 **Solution:** Replace the locked inode with a fresh readable one. `mv` is atomic; mode bits live in the inode,
     not the directory entry. Two flavors cover every case we hit on a real machine:
1598
1599 **Flavor A  locked file is mode `000` (owner can't even READ it):**
1600 NOPASSWD sudoers (`/etc/sudoers.d/nullexit`) already includes `/usr/bin/python3`. So we read via `sudo python3`
      (which has the DAC override root has) and pipe the bytes around as base64 in user space, where the
     redirect happens with the user's normal umask = `0644`:
1601 ```bash
1602 sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).
     decode())" > /tmp/snap.b64
1603 base64 -d < /tmp/snap.b64 > <FILE>.new
1604 mv -f <FILE>.new <FILE>
1605 ls -la <FILE>          # mode is now 0644
1606 ```
1607 The old `mode=000` inode is unlinked by `mv`; the new `mode=0644` inode (created by base64-decode via the
     calling user's umask) takes the path. No chmod ever called. The parent directory just needs to be writable
     by you (it always is, since you own it).
1608
1609 **Flavor B  locked file is mode `0444` (you can READ but cannot WRITE; `scripts/sync-rules.py` needs to re-
     write):**
1610 You own the file and can read it, you just cannot truncate/overwrite it. The simplest path is to rename it
     aside and let the writer create a fresh inode from scratch (which gets the umask-default `0644`):
1611 ```bash
1612 mv <FILE> /tmp/old.<FILE>.$$
1613 python3 scripts/sync-rules.py    # now writes <FILE> cleanly at mode 0644
1614 ```
1615 Or apply Flavor A if you'd prefer to preserve the contents while resetting the mode.
1616
1617 **Why this is the canonical recovery, not a hack:**
1618 - Mode bits live in the inode, not the path. `mv -f` replaces the path's inode reference; the old locked inode
     drops out of the directory entirely (dentry eviction) and loses its last link, so the kernel reclaims it.
     The new inode (whatever it is: a freshly-written one, or a base64-decoded copy) gets the path.
1619 - The only prerequisite to apply Flavor B's `mv` is: you own the parent directory (so you can unlink entries
     from it) AND you have a NOPASSWD-capable binary that can read the byte stream if mode prohibits user read.
1620 - This pattern generalizes. Any time a script ends up producing a "locked file that the next run can't update"
     situation, the same trick applies without a single chmod anywhere.
1621
1622 **Empirical verification (user machine, July 1, 2026):**
1623 - `.env` was `mode=000` (owner $USER, i.e. you): Flavor A unblocked it in 3 commands, no chmod.
1624   - Before: `----------  1 $USER  staff  899 Jun 28 07:07 .env`
1625   - After:  `-rw-r--r--@ 1 $USER  staff  899 Jul  1 20:39 .env`
1626   - `docker compose config` immediately picked up `TS_AUTHKEY` + `WIREGUARD_PRIVATE_KEY` substitution.
1627 - `compiled_rules.txt`, `ip_blocklist.ipset`, `data/filters/1782645604.txt` were `mode=0444`: Flavor B (`mv`-
     aside) unblocked all three.
1628   - scripts/sync-rules.py then ran clean: 279587 block rules + 171 allow rules + 16710 IP entries compiled; new
      files came out at `0644`.
1629 - `toggle.sh` then went end-to-end in 48s (warp / routing-fix / adguardhome / tailscale all healthy, gateway
     mesh-joined, DNS hijack verified single-server).
1630
1631 **When this trick is appropriate vs. inappropriate:**
1632 -  You own the parent directory and the file (typical desktop scenario).
1633 -  The lock is a leftover from a removed `chmod` call that you want off disk forever.
1634 -  NOPASSWD sudo includes at least one interpreter (`python3` on this system) that can be used to read mode-0
     files. If it does not, add it first: `<USER> ALL=(ALL) NOPASSWD: /usr/bin/python3`.
1635 -  The file is owned by another user (root, another account) and the lock is intentional  respect it; do not
     bypass another user's security boundary.
1636 -  The parent directory is owned/writable only by another user  escalate properly via sudo, or stop and ask the
      user.
1637 -  You do not actually own the file and there is a legitimate reason for its lock state  restore the file via a
     trusted source rather than circumventing the perms.
1638
1639 **Reusable cookbook:**
1640
1641 ```bash
1642 # Quick diagnostic: figure out which flavor you need.
1643 WHOAMI=$(whoami)
1644 ls -la <FILE>
1645 stat -f 'mode=%Sp owner=%Su group=%Sg' <FILE>
1646 # mode=--------- or mode=----r----- : you cannot even read it  Flavor A.
1647 # mode=r--r--r-- but writer fails  : you are blocked from write but can read  Flavor B.
1648 # mode=rw-r--r--                   : not actually locked at all; unrelated write failure.
1649
```

```
# Flavor A (mode=000):
sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).decode())" > /tmp/snap.b64
base64 -d < /tmp/snap.b64 > <FILE>.new && mv -f <FILE>.new <FILE> && rm -f /tmp/snap.b64

# Flavor B (mode=0444, file is readable):
mv <FILE> /tmp/old.<FILE>.$(date +%s)
# (let the original writer recreate <FILE>; new file lands at mode 0644)

# Verify after either flavor:
ls -la <FILE>             # mode should be 0644
<your normal validation command>
```

15.10 Tor Container

15.10.1 July 11, 2026: Tor Container Hardening & obfs4proxy Restoration
Symptom:
1. The Tor container entered a fatal crash loop upon boot, throwing `exec: "obfs4proxy": executable file not found in $PATH` when `TOR_USE_BRIDGES=true`.
2. When the user pasted Tor bridge lines containing comments (`#`) or blank lines into `tor-bridges.txt`, the container failed to validate the file properly and crashed.

Root Cause:
- **Bug 1 (Missing Pluggable Transports):** The Tor Dockerfile was built on `debian:bullseye-slim`, which had silently dropped or failed to resolve the `obfs4proxy` package in its latest apt repositories.
- **Bug 2 (Fragile Validation):** The `entrypoint.sh` bridge validation check merely checked `[ -s tor-bridges.txt ]` (file size > 0). It did not verify if the file actually contained valid, uncommented bridge lines. Passing empty lines or comments tricked the validation script into appending invalid `Bridge` directives to the `torrc`, causing the Tor daemon to instantly exit.

Resolution:
- **Upgraded Base Image:** Shifted the Tor Dockerfile base image to `debian:bookworm-slim`, restoring access to the `obfs4proxy` pluggable transport package.
- **Robust Bridge Validation:** Replaced the fragile `[ -s ]` check with a strict `grep -vE '^\s*(#|$)'` command that strips comments and empty lines. If no valid lines remain, the container safely falls back to standard public guard relays instead of crashing, ensuring Tor remains available even with a misconfigured bridge file.
- **Image Pinning:** Explicitly pinned `image: nullexit-tor:v1.0.0` in `docker-compose.yml` to prevent arbitrary local builds from drifting to `:latest`.

15.10.2 July 12, 2026: Tor ControlPort Dynamic Password Authentication & User Config
Symptom: Querying the Tor Control Port (e.g., via the `/sweep` script or `check_circuits.py`) returned empty responses or connection errors. Logs showed that Tor automatically closed its ControlPort:
`You have a ControlPort set to accept unauthenticated connections from a non-local address. ... That's so bad that I'm closing your ControlPort for you.`

Root Cause: Docker port mapping (especially under Colima on macOS) requires container sockets to bind to `0.0.0.0` to be reachable from the host. However, when Tor's ControlPort is bound to `0.0.0.0` (non-local) with no authentication, Tor closes it by default as a safety precaution. Sourcing SOCKS5/ControlPort to `127.0.0.1` inside the container was not an option as it broke the Docker bridge routing path.

Resolution:
- **Dynamic Password Authentication**: Updated the Tor container's `entrypoint.sh` to read `TOR_PASSWORD` and dynamically generate its HMAC hash via `tor --hash-password`, adding `HashedControlPassword <hash>` to `torrc`. This secures the ControlPort on `0.0.0.0` so Tor keeps it open.
- **Unified Credentials in `.env`**: Exposed `ADGUARD_USER=admin` and `ADGUARD_PASSWORD=nullexit` in `.env` to serve as the unified gateway credentials. Sourced `TOR_PASSWORD` directly from `ADGUARD_PASSWORD` in `docker-compose.yml`.
- **Client Authentication**: Updated `scripts/check_circuits.py` to fetch `TOR_PASSWORD` (defaulting to `ADGUARD_PASSWORD` or `nullexit`) and authenticate with the ControlPort using `AUTHENTICATE "<password>"` to regain access.

> See also 11.3 (Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification) for the design-level fail-closed guarantee.

15.11 macOS Platform Quirks

15.11.1 macOS Application Firewall Silently Blocking AirDrop & Continuity
Problem: Users of complex network extensions (Tailscale, LuLu, Docker, etc.) often run into an issue where AirDrop, AirPlay, and Universal Clipboard silently fail, even with the gateway inactive and Wi-Fi/Bluetooth correctly configured. The internal Apple Wireless Direct Link (`awdl0`) interface appears healthy, but connections never establish.

Root Cause: The built-in macOS Application Firewall has a specific list of explicitly blocked applications. Apple's own core networking services (e.g., `/usr/libexec/sharingd`, `/usr/libexec/rapportd`, `/usr/libexec/avconferenced`) can be silently added to the "Block incoming connections" listeither through accidental "Deny" clicks on permission prompts, execution of aggressive privacy-hardening scripts, or a known macOS bug that retains strict blocks even after toggling the "Block all incoming connections" setting off. Since these are background daemons, macOS provides no visible error.

Fix: The firewall rules must be cleared via the terminal (or System Settings > Network > Firewall > Options):

```bash
sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/sharingd
sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/rapportd
```

Once unblocked, `sharingd` immediately resumes broadcasting mDNS/Bonjour discovery packets, and AirDrop restores instantly without requiring a reboot.

15.11.2 Standalone Tailscale Daemon Freeze after macOS Sleep/Wake

Problem: When the macOS host goes to sleep and wakes up, the network interfaces and routing tables are rebuilt. The standalone `tailscaled` daemon (installed via Homebrew) occasionally fails to handle this transition and becomes completely frozen/unresponsive. In this state, any command like `tailscale down` or `tailscale status` hangs indefinitely, and all DNS requests sent to the local Tailscale resolver (`100.100.21.8`) time out, causing a total internet blackout on the host.

Fix: Enhanced the Tailscale verification and teardown logic in `toggle.sh` and `recover.sh`:

1. **Unresponsive Daemon Detection:** Instead of assuming the daemon is healthy just because the Unix socket `/var/run/tailscaled.socket` exists, the scripts now actively run a status check with a 5-second timeout (`run_with_timeout 5 tailscale status`). If the check times out (exit code 143), the daemon is classified as unresponsive.
2. **Auto-Recovery Loop:** If the daemon is unresponsive, the script now automatically attempts to restart the service using `brew services restart tailscale` (falling back to `sudo -n brew services restart tailscale`).
3. **Graceful Retry:** Once restarted, the script pauses for 3 seconds for the daemon to initialize before proceeding with connection or teardown commands.

15.11.3 macOS Sharing Services (AirDrop/AirPlay) Freeze on Network Transitions

Problem: When the gateway is turned on or off, the scripts perform network configuration cleanups which include flushing the routing tables (`route -n flush`), restarting the `mDNSResponder` daemon (`killall -HUP mDNSResponder`), and power-cycling the Wi-Fi interface (`setairportpower` or `ifconfig en0 down/up`). While AirDrop BLE discovery continues to function, the actual file transfers stall at "Waiting..." indefinitely.

Root Cause: The rapid tear-down and rebuild of the local routing table and Wi-Fi interface causes Apple's core sharing and connection daemons (`sharingd` and `rapportd`) to lose their socket bindings to the mDNS/Bonjour discovery interface. Instead of reconnecting gracefully, they enter a wedged state and try to route peer-to-peer (AWDL) IPv6/TCP transfer traffic into the default gateway (the Tailscale interface `utun`) instead of the direct wireless link, causing the TLS verification handshake to time out.

Fix: Integrated an automatic sharing services reset into both `toggle.sh` and `recover.sh`:

1. The script now automatically runs `sudo -n killall sharingd rapportd` at the end of both the **START** path and the **STOP** path (as well as inside `recover.sh`).
2. Killing these daemons forces macOS to immediately relaunch them, binding them fresh to the newly initialized network interfaces and routing tables, allowing AirDrop to work seamlessly while the gateway is active.

15.11.4 Chrome Remote Desktop Connection Failures (Tailscale on Remote Device)

Context: When attempting to access a remote machine via Chrome Remote Desktop (CRD), the connection fails or hangs if Tailscale is active **on the remote device**. The connection works fine even if Tailscale (and the exit node) is active **on the local client/viewer device**, but only if Tailscale is disabled on the remote machine. *(This is distinct from the CRD-through-WARP latency limitation in 14.3.)*

Why: Having Tailscale active on the remote machine creates virtual network interfaces and DNS/routing modifications that conflict with Chrome Remote Desktop's WebRTC bindings and signaling traffic.

Workaround:

- **Use Tailscale Native RDP/RustDesk (Recommended):** Connect directly to the remote device's Tailscale IP (`100.X.Y.Z`) via an RDP or RustDesk client. This is faster and avoids third-party relay conflicts.
- **Disable Tailscale on the Remote Device:** If you must use CRD, disable Tailscale on the remote machine before connecting.

15.11.5 tailscaled Daemon Broken State (brew services)

Symptom: `brew services list` shows `started` but `tailscale status` shows "stopped".

Fix: `sudo brew services restart tailscale`. Toggle script now detects and auto-restarts.

15.11.6 Apple Silently Removed the `airport` Binary in macOS Sonoma 14.4 (March 2024)

What happened: `watcher.sh`'s `detect_lan_p2p_mode()` was written to use the `airport` CLI tool at:

```
/System/Library/PrivateFrameworks/Apple80211.framework/Versions/Current/Resources/airport
```

This path returned `exit 127: no such file or directory`. The function silently fell back to `reason="no-wifi"` even while the Mac was actively connected to WPA2 Enterprise Wi-Fi.

Why Apple removed it: `airport` was never a real public binary it lived inside a **private framework** and was always technically unsupported. Apple deprecated it quietly years ago and finally removed it in **macOS Sonoma 14.4 (March 2024)**. The removal is part of a broader privacy clampdown on Wi-Fi metadata:

- Starting in **macOS 14.5**, BSSIDs and other Wi-Fi association details are **redacted** (`<redacted>`) in most tools, including the official replacement `wdutil`
- Apple's official recommendation is to use the **CoreWLAN framework** (Swift/Objective-C) with proper entitlements useless for a bash script

- The official CLI replacement `wdutil` requires `sudo` for most useful output also useless for a background LaunchAgent

Fix: Replaced `airport -I | awk '/link auth/'` with:
```bash
system_profiler SPAirPortDataType | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF ; exit}'
```
This returns human-readable strings like `WPA2 Enterprise`, `WPA2 Personal`, or `None` for the currently associated network without `sudo`, and without any removed private binary.

Confirmed working output on this machine:
```
P2P detect: security='Enterprise'  allow=false (reason: 802.1x-enterprise)
```

> ** Risk: `system_profiler` may not survive forever either.** `system_profiler SPAirPortDataType` still works as of macOS Sequoia (15.x) without sudo and without redaction. However, given Apple's trajectory on Wi-Fi privacy, this could be locked down in a future release. If `security` comes back empty on a future macOS version while the Mac is actively on Wi-Fi, the function will silently fall back to `allow=false` (safe default), but the detection will be broken again. The long-term fix would be a small Swift helper using CoreWLAN that we compile once and bundle in the repo. *(This forward-looking caveat is also tracked as an observation in 16.)*

15.11.7 Changing macOS Local Hostname
Context: Users may wish to change their Mac's computer name/hostname in macOS System Settings.
Effect: Changing the Mac's hostname will **not** break Nullexit (which relies on internal routing/localhost). However, Tailscale will automatically detect this change and update the machine name on the Tailnet.
- The underlying Tailscale `100.x.x.x` IPv4 address will remain exactly the same.
- The **MagicDNS name** will change to match the new hostname. Users must remember to update their SFTP/SSH clients to use the new MagicDNS address.

15.11.8 Shell-Language Landmines: macOS `/bin/bash` 3.2 + `set -e` (Field Notes)

This project's lifecycle scripts (`toggle.sh`, `recover.sh`, `common.sh`,) run under `#!/bin/bash`, which on macOS is **GNU bash 3.2.57 (2007)** Apple has never shipped a newer bash because bash 4+ is GPLv3. (Homebrew's bash 5 lives at `/opt/homebrew/bin/bash` but is *not* the shebang target, and cannot be assumed on PATH a bare `bash --version` on a stock Mac reports 3.2.) Every script here also runs `set -e` (exit-on-error) with an `ERR`/`INT`/`TERM`/`HUP` trap. That combination **ancient bash + `set -e` + traps + `set -x` xtrace redirection** is a minefield. The traps below were all hit and fixed during development; document them so they are not re-discovered.

A. `set -e` + `if  then  else  fi` can abort at the `else` (bash 3.2 heisenbug).
The nastiest one. A construct as innocent as:
```bash
if [ -f "$MARKER_FILE" ]; then
  X="yes"
else
  X="no"
fi
```
aborted `toggle.sh` at init under `set -e` whenever the file was **absent** the false status of the `[ -f  ]` test leaked through the compound and `set -e` killed the script (`EXIT: code=1 phase=init`), before any real work ran. Two things made this vicious: (1) it is intermittent/context-dependent it did **not** reproduce in isolated `bash -c ''` snippets of the identical block; it only manifested inside the full script with the `set -x` xtrace fd-redirect (`exec 2> >(tee )`) active. We only proved causation by **bisecting the live script** (replace the block with a trivial assignment the script sailed past init). (2) Static tools (`bash -n`, `shellcheck`) cannot see it it is a runtime interaction. **Safe pattern:** default-assign, then an `if` with **no `else`** (an `if` whose false condition has no else-branch yields status 0):
```bash
X="no"
if [ -f "$MARKER_FILE" ]; then X="yes"; fi
```
Confirmed root cause (2026-07-13, reproduced head-to-head). The trigger is specifically **`set -x` plus a `PS4` that contains a command substitution** here the `DEBUG_TRACE` prompt `PS4='+ [$(date )] '`. Forcing the false-condition/else path (`gateway_state` `stopped`, so `if [ "$(gateway_state)" = "running" ]` takes the else/START branch) on stock `/bin/bash` 3.2.57: with the **default `PS4`** it survives (3/3, exit 0); with toggle's **`$(date)` `PS4`** the `ERR` trap fires on the else branch (`rc=1` from the false condition, 3/3, exit 42). The `$(date)` runs a subshell on *every* traced command; while tracing the `if`-condition it perturbs `$?` so the condition's non-zero status leaks past the normal `if`-exemption and trips the `ERR` trap. This is a true **heisenbug** enabling tracing (`DEBUG_TRACE=true`) is what *causes* the abort, which is exactly why it's invisible with `DEBUG_TRACE=false` (the default) and why every isolated repro on the *default* `PS4` passes (this misled an entire debugging session into briefly "disproving" the landmine). It aborted the gateway START decision (`toggle.sh:689`) under `DEBUG_TRACE=true`; with `DEBUG_TRACE=false` the gateway starts cleanly (verified: 192 s, all five containers healthy). **Fixed (2026-07-14).** `DEBUG_TRACE`'s `PS4` now uses the subshell-free `${SECONDS}` (elapsed seconds since the shell started) instead of `$(date "+%H:%M:%S")`, removing the per-traced-command subshell that perturbed `$?`. bash 3.2 has no subshell-free *wall-clock* token for `PS4` (`$EPOCHSECONDS` is 5.0+, `printf '%()T'` is 4.2+), so the per-line stamp is now elapsed-seconds the absolute start time is still logged once in the `DEBUG_TRACE enabled ` header line, so wall-clock is recoverable. Bonus: this also drops a `date` fork on *every*

traced line (previously thousands). **Verified head-to-head on stock `/bin/bash` 3.2.57:** `toggle.sh --restart` with `DEBUG_TRACE=true` now completes `code=0` with tracing active straight through the exact decision that aborted before the trace shows `gateway_state  stopped` (`common.sh:229`) feeding the false `if [ "$(gateway_state)" = "running" ]`, then the else branch reaching `Action: STARTUP GATEWAY` (1436 traced lines, no `ERR` trap firing, all five containers healthy). `DEBUG_TRACE` is now safe to use on the STOP/START decision path.

More generally, any command whose "failure" is normal control flow (`grep -q`, `[ ]`, `diff`, `kill -0`) that ends up as the *last* command of a function/branch/script under `set -e` is a landmine; guard it (`|| true`) or restructure so a non-zero status can't be the tail value.

B. `BASH_XTRACEFD` and the `exec {var}>>file` dynamic-fd form are bash 4.1+ (hard-fail on 3.2).
Our `DEBUG_TRACE` feature wants xtrace (`set -x`) sent to `output.log` while keeping the terminal clean the clean way is `exec {BASH_XTRACEFD}>>"$LOG_FILE"`. On 3.2 this is a **syntax/runtime error**: `exec: {BASH_XTRACEFD}: not found` (dynamic-fd allocation `{var}>>` didn't exist until 4.1, and `BASH_XTRACEFD` itself is 4.1+; on 3.2 it's just an ordinary variable that `set -x` ignores). Left unguarded it would abort the script at startup i.e. the debug feature would break the very thing it's meant to observe. **Safe pattern:** branch on `${BASH_VERSINFO[0]}.${BASH_VERSINFO[1]}` on 4.1 use `exec 9>>"$LOG_FILE"; export BASH_XTRACEFD=9`; on 3.2 fall back to `exec 2> >(tee -a "$LOG_FILE" >&2)` (xtrace goes to stderr, which we duplicate to the log; the terminal also shows it, which is acceptable). Never use the `exec {var}>>` form unless you *know* you're on 4.1.

C. `trap  ERR` + `$LINENO` misattributes the failing line on 3.2.
When a command fails under `set -e`, our trap runs `cleanup_handler ERR $LINENO`. On bash 3.2, `$LINENO` inside an `ERR` trap frequently reports the line of the **enclosing `fi`/closing brace of a compound statement**, not the true failing line. During the STOP/START-decision bug this produced `Script failed at line 1202` where 1202 is a bare `fi` actively misleading. **Mitigation:** don't trust the trap's line number as gospel on macOS; corroborate with the `EXEC:`/`EXIT:` lifecycle breadcrumbs and the `DEBUG_TRACE` xtrace (which prints `file:line:function` via a custom `PS4`), which pin the *actual* last-executed command.

D. `set -e` is not inherited by command substitutions / subshells the way people expect, and is suppressed inside `if`/`&&`/`||` conditions.
This cuts both ways and is a frequent source of "why didn't it fail?" and "why *did* it fail?": a non-zero command used as an `if` condition (or left/right of `&&`/`||`) is *exempt* from `set -e` (by design), but the moment that same call is a bare statement it aborts. `is_gateway_active()` relies on this (it's always called as `if is_gateway_active; then `), which is precisely why moving the decision to a marker file instead of a probe that must be `if`-guarded to be safe is more robust (see 15.12.19). Also note `local X=$(cmd)` masks `cmd`'s exit status (the `local` builtin's status wins) a `shellcheck` SC2155 we see repeatedly; harmless when we don't check `$?`, but a real footgun if you do.

E. Process substitution (`>()`, `<()`) is a bashism, and its lifecycle is loose.
Our 3.2 xtrace fallback (`exec 2> >(tee -a "$LOG_FILE" >&2)`) uses process substitution fine under bash, but (1) it is **not** POSIX `sh`, so these scripts must never be invoked as `sh script.sh`; and (2) the background `tee` it spawns is reaped asynchronously, so the very last few lines written just before exit can occasionally be truncated in the log. Acceptable for a debug trace; do not rely on it for critical last-gasp logging (use a direct `>> "$LOG_FILE"` append for must-capture lines, as the `EXIT` breadcrumb does).

F. GNU-vs-BSD userland divergence (not bash itself, but same class of pain).
macOS ships BSD variants of the core tools, so idioms differ from Linux: `sed -i ''` (BSD requires the empty backup-suffix arg) vs `sed -i` (GNU); `stat -f%z` (BSD) vs `stat -c%s` (GNU); `date` format flags; `jot` (BSD) vs `shuf` (GNU) for random port generation; `route`/`networksetup`/`dscacheutil` (macOS-only) vs `ip`/`resolvectl` (Linux). This is *why* the codebase branches on `[[ "$OSTYPE" == "darwin"* ]]` everywhere `toggle.sh` and `recover.sh` are each single cross-platform scripts that dispatch on `$OSTYPE` (an early Linux block runs and exits before the macOS body). `timeout(1)` doesn't exist on stock macOS at all hence the pure-bash `run_with_timeout` in `common.sh`.

Working rules distilled from the above:
- Assume **bash 3.2** semantics for anything under `#!/bin/bash` on macOS; gate any 4.0+ feature (`BASH_XTRACEFD`, dynamic `{var}` fds, associative arrays, `${var^^}` case-conversion, `mapfile`/`readarray`, negative array indices) behind a `BASH_VERSINFO` check or avoid it.
- Under `set -e`, never let a "benign non-zero" command (`[ ]`, `grep -q`, `kill -0`, `diff`) be the tail statement of a branch/function/script; use no-`else` `if`s, `|| true`, or explicit `$?` handling.
- Prefer **persisted state (marker files)** over **live probes** for control-flow decisions a probe forces an `if`-guard and can hang/abort; a file read cannot (see 15.12.19).
- `bash -n` and `shellcheck` catch syntax and many smells but **cannot** catch `set -e`/trap runtime interactions a **live run is mandatory** to validate lifecycle-script changes; the `DEBUG_TRACE` + breadcrumb instrumentation exists precisely to make those live runs diagnosable.
- Never run these scripts as `sh` they are bashisms end to end.

15.12 Misc Resolved

15.12.1 Gluetun Resets nftables FORWARD Rules
Symptom: FORWARD RELATED,ESTABLISHED rule disappears after Gluetun health check.

Fix: `scripts/routing-fix.sh` re-adds to both iptables backends every 30 seconds. Legacy backend (policy ACCEPT) acts as safety net.

15.12.2 Native SSH & SFTP Security (Zero Local Attack Surface)
Context: macOS "Remote Login" (`sshd`) and "File Sharing" (`smbd`) open ports on the local network (e.g. `172.x.x.x`), exposing the host to brute-force attacks on public Wi-Fi.

Implementation: The script strictly enforces `tailscale up --ssh=true` whenever the mesh state is reset. This empowers the user to completely disable native macOS Remote Login and File Sharing. This provides an incredible zero-trust security advantage: even if an attacker steals your Mac username and password, they **cannot** access your files remotely. Disabling the native services completely closes the listening ports on the local Wi-Fi interface (`172.x.x.x`), rendering the Mac inaccessible to local network logins. Meanwhile, Tailscale intercepts port 22 traffic exclusively over the encrypted mesh and authenticates cryptographically based on your Tailscale SSO identity. Stolen Mac passwords are mathematically useless without first bypassing your Tailscale Two-Factor Authentication and registering a device onto the mesh.

Cross-Platform File Exchange: To easily browse and exchange files securely, simply use an SFTP client and connect to your Mac's MagicDNS name on port 22. Recommended clients: **WinSCP** or **FileZilla** (Windows), **Cyberduck** (macOS), and **FE File Explorer** or **Solid Explorer** (iOS/Android).

15.12.3 Logging Architecture
For debugging, logs are strictly segmented based on the component's lifecycle:
- **`output.log` (Host-side):** Contains all standard error (`stderr`) and verbose output from the host scripts (`toggle.sh`, `recover.sh`, `setup.sh`). Since the `rule-compiler` container is ephemeral and deleted after running, its logs (and any Python errors from `scripts/sync-rules.py`) are extracted and appended to this file before deletion.
  - **Log Rotation Policy:** On startup, `toggle.sh` checks if `output.log` exceeds **1GB** (`1,073,741,824` bytes). If it does, it performs a metadata-only rename (`mv output.log output.log.old`), preserving history while resetting the active log to 0 bytes. This rename-based rotation avoids the heavy disk I/O and CPU overhead of rewriting a massive file line-by-line. The threshold is intentionally large because with `DEBUG_TRACE=true` the log accumulates a full command-by-command execution trace (effectively the framework's entire compilation/warning history), which is valuable to retain for post-mortems; a smaller cap would rotate that history away too quickly.
  - **Egress Probe Logging:** The background host-egress leak prober checks if stdout is a TTY (`[ -t 1 ]`) and will only print live status updates (using carriage return `\r`) or duplicate console warnings in interactive mode. It automatically detects macOS/BSD vs. Linux environments to dynamically construct timestamps with date and time (omitting milliseconds on macOS to prevent high CPU utilization from spawning sub-second processes).
- **`docker logs <container>` (Guest-side):** The persistent containers (`warp`, `tailscale`, `routing-fix`, `adguardhome`) use the Docker `json-file` logging driver with a strict `max-size` (1m-10m) to prevent VM disk exhaustion. Use standard `docker logs` to view them.

15.12.4 Pre-Flight Check 1: Wrong `tailscale ping` Flag
Problem: Check 1 of the pre-flight connectivity checks used `--c 1` (double dash) but `tailscale ping` accepts `-c 1` (single dash). The invalid flag caused the command to exit with code 1, marking the check as FAIL even though the gateway was reachable.

Fix: Changed `--c 1` to `-c 1`.

15.12.5 Pre-Flight Check 1: DERP Relay Exit Code
Problem: Even with the correct `-c 1` flag, `tailscale ping` exits with code 1 when the connection goes through a DERP relay (`direct connection not established`) even though a pong was successfully received (28ms via DERP(tor)). The check treated relayed connections as failures.

Fix: Changed the check from relying on exit code to piping stdout through `grep -q "pong"`. This accepts relayed connections as valid (they work fine for the exit node use case), while still failing on true timeouts or unreachable peers.

15.12.6 IPv6 Exit-Node Leak
Problem: Tailscale supports IPv6, but WARP/gluetun is IPv4-only. IPv6 packets forwarded through the exit node would leak through the Docker bridge unencrypted.

Fix: Added `ip6tables -A FORWARD -i tailscale0 -j DROP` to `post-rules.txt`, blocking all IPv6 forwarded traffic from the Tailscale interface. *(This is the fix referenced by the IPv6-leak scenario in 15.6.1.)*

15.12.7 SOCKS5 Proxy Threading Race Condition
Problem: The `forward` function in `scripts/socks5-proxy.py` called `select.select([src, dst], [], [])` which raised `ValueError: file descriptor cannot be a negative integer (-1)` when one thread closed a socket while the other thread was selecting on it.

Fix: Added `ValueError` exception handling around `select.select()` and `recv()` calls. When a file descriptor becomes negative, the thread breaks out of the forwarding loop cleanly instead of crashing.

15.12.8 Docker Healthcheck: Multi-line Python in CMD Array
Problem: The socks-proxy healthcheck used a `CMD` array with a multi-line Python string. YAML collapsed the newlines, causing the `#` comment to comment out the rest of the code and the `assert` statement to be mangled into `as`.

Fix: Changed from `["CMD", "python3", "-c", "..."]` to `["CMD-SHELL", "python3 -c '...'"]` with a single-line Python expression. Uses a real SOCKS5 handshake (sends `[5, 1, 0]`, expects `[5, 0]`) instead of a simple TCP port check.

15.12.9 Graceful Shutdowns Leave DNS Hijacked
Problem: The gateway overrides the macOS host's DNS settings (via `networksetup`) to route queries to the AdGuard container. Because macOS persists `networksetup` changes to disk, shutting down the Mac while the gateway is running causes the Mac to boot up later with hijacked DNS. Since the Colima VM and Docker containers don't auto-start on boot, the user would have no internet access until they manually run the toggle script.

```
1860
1861 **Root Cause:** The `toggle.sh` script exits after the gateway starts. There was no background daemon listening
       for macOS system shutdown signals to revert the network settings.
1862
1863 **Fix:** Wrapped the background `caffeinate` process (used to prevent system sleep) in a bash `trap` command.
      The trap listens for `SIGTERM`, `SIGINT`, and `SIGHUP`. When the user initiates a graceful shutdown, macOS
      sends `SIGTERM` to the background trap, which instantly executes `recover.sh` to flush the host DNS back to
       `1.1.1.1` right before the machine powers off.
1864 *Note:* We explicitly use `recover.sh` instead of `toggle.sh` because during a shutdown, the OS limits process
      cleanup time and is already independently sending termination signals to Docker and Colima. Trying to run `
      docker compose down` inside a shutdown trap creates a race condition that can hang the script until macOS
      forcefully `SIGKILL`s it. `recover.sh` abandons container management and surgically flushes the macOS
      network settings in milliseconds, guaranteeing completion before the OS pulls the plug.
1865
1866 #### 15.12.10 Linux Native Wi-Fi Bouncing
1867 **Problem:** The generated `scripts/recover-linux.sh` incorrectly retained the macOS `networksetup` commands,
      which would crash and fail to bounce Wi-Fi on Linux hosts.
1868
1869 **Fix:** Swapped the macOS logic for native Linux commands. The script now attempts `nmcli radio wifi off/on`
      first, and gracefully falls back to `ip link set wl... down/up` if NetworkManager is unavailable.
1870
1871 #### 15.12.11 TS_AUTH_ONCE Cloud Footgun & Ephemeral Keys
1872 **Decision:** The compose file uses `TS_AUTH_ONCE=true` to prevent authentication loops. However, on headless
      cloud VPS deployments, if a standard auth key expires (usually 90 days), the container will silently drop
      off the mesh. We documented this trade-off explicitly in `docker-compose.yml` and recommended the use of **
      Tailscale Ephemeral Keys** for cloud deployments, which automatically clean themselves up and bypass this
      issue.
1873
1874 #### 15.12.12 Hardcoded AdGuard Credentials
1875 **Decision:** AdGuard is deployed with the hardcoded credentials `admin / nullexit`. This is perfectly safe
      behind a NAT on a local macOS laptop. However, if deployed on a cloud host with exposed ports, this becomes
       a severe security risk. We added a stark warning comment in `docker-compose.yml` to ensure cloud users
      change this before deployment.
1876
1877 #### 15.12.13 Routing Fix: 30-second polling vs Netlink Socket
1878 **Decision:** Instead of building a complex Netlink socket listener (`ip monitor`) to react instantly to
      iptables and routing flushes from Gluetun, we retained the 5-second `sleep` loop in `scripts/routing-fix.sh
      `.
1879 - **Reasoning:** Building a netlink listener in pure bash on Alpine is incredibly brittle and complex (
      especially for iptables events). The 30-second polling loop is bulletproof. The only trade-off is a
      potential 1-4 second stutter if Gluetun reconnects, which was deemed an acceptable edge case for absolute
      reliability.
1880
1881 #### 15.12.14 Cache Poisoning and URL Rot in scripts/sync-rules.py
1882 **Problem:** If a remote blocklist URL went permanently offline (404 Not Found), the Python `urllib` request
      would return the 404 HTML string. The script would aggressively regex-parse this HTML, find no domains, and
       overwrite the healthy local cache with a 0-domain file. This effectively disabled ad-blocking until the
      URL was fixed.
1883
1884 **Fix:** Implemented a defensive programming sanity check in `scripts/sync-rules.py`. Before overwriting the
      cache, the script now verifies that the compiled domain count is greater than 10 (and 1 for IP feeds). If
      it falls below this threshold (indicating a catastrophic 404 failure across multiple URLs), it aborts the
      cache overwrite, raises a `ValueError`, and keeps the previous day's healthy cache intact.
1885
1886 #### 15.12.15 Gluetun nf_tables Parser Crash (Bash Interpolation Failure)
1887 **Problem:** Attempting to make the TCP MSS dynamically configurable by using a standard bash variable inside `
      post-rules.txt` (`iptables ... --set-mss ${GATEWAY_MSS:-1120}`) caused the Gluetun container to
      catastrophically crash during startup. Gluetun's internal parser executes `post-rules.txt` line by line
      without a shell, meaning the variable was never expanded. It literally attempted to inject the string `${
      GATEWAY_MSS:-1120}` into iptables, resulting in a fatal `bad value for option "--set-mss"` error.
1888
1889 **Fix:** Removed the bash variable from `post-rules.txt` and reverted it to a pure hardcoded integer. To retain
       dynamic `.env` configuration, we moved the interpolation logic to `toggle.sh`. Right before Docker starts,
       the host script parses `GATEWAY_MSS` from the `.env` file and uses a cross-platform `sed` command to
      dynamically overwrite the integer directly inside `post-rules.txt`. This perfectly mimics manual
      configuration and completely bypasses Gluetun's strict parser.
1890
1891 #### 15.12.16 AdGuard Filter Redundancy & Memory Optimization
1892 **Problem:** AdGuard Home does not perform cross-list deduplication in memory. When it loads its own native
      subscription filters (e.g., `AdGuard DNS filter`) alongside our massive `compiled_rules.txt`, overlapping
      rules are loaded twice, wasting significant RAM in the Colima VM. The UI would show ~500k rules,
      representing the sum of all lists rather than unique domains.
1893
1894 **Fix:** Updated `scripts/sync-rules.py` to intelligently cross-reference AdGuard's configuration and
      deduplicate our list against it.
1895 1. The script now reads `adguard/conf/AdGuardHome.yaml` to dynamically fetch the exact URLs of any enabled
      native AdGuard filters.
1896 2. The parsing engine was upgraded to normalize basic AdGuard syntax (`||domain^`) back into raw base domains
      during the build process.
1897 3. The script subtracts the native AdGuard domains from our custom blocklist *before* compiling, immediately
      purging ~83,000 completely redundant rules.
```

```
4. The normalized base domains then pass through our subdomain optimizer, which squashes an additional ~50% of the remaining rules.

This reduces the final compiled output from ~325k to ~281k rules on the `heavy` profile, saving ~45k rules from being loaded into memory twice and reducing the overall RAM footprint of the `adguardhome` container.

#### 15.12.17 Kernel-Level IP Blocklist (Threat Intelligence Firewall)
**Problem:** DNS sinkholing cannot stop malware or botnets that bypass DNS entirely by hardcoding direct IP addresses (a common technique for C2 communication). These connections pass straight through AdGuard unseen.

**Fix:** Added a second compilation pipeline to `scripts/sync-rules.py` that runs on every startup alongside the DNS pipeline.
1. The compiler concurrently fetches four curated threat-intelligence feeds: **Feodo Tracker** (abuse.ch botnet C2 IPs), **Spamhaus DROP** (IPs allocated to criminal organizations), **Emerging Threats** (Proofpoint  active C2 and scanners), and **CINS** (active brute-force sources).
2. All entries are normalized via Python's `ipaddress` module. Any IP or CIDR that overlaps with RFC1918 private ranges, loopback, link-local, or the Tailscale CGNAT range (`100.64.0.0/10`) is stripped to prevent accidentally locking users out of their own LAN or mesh.
3. The cleaned list (16,721 unique IPs/CIDRs) is written as an `ipset restore` file using an **atomic swap pattern**: `create_new  populate  swap with live  destroy_new`. This guarantees the live `block_malicious` ipset is never empty during a reload.
4. `scripts/routing-fix.sh` watches the file's mtime every 30 seconds. On change it runs `ipset restore` and idempotently re-injects `FORWARD DROP` rules for both `src` and `dst` into both iptables backends (nftables + legacy).
5. `docker-compose.yml` mounts `adguard/work/userfilters/` into `routing-fix` as `/userfilters:ro` and adds `rule-compiler: service_completed_successfully` to its `depends_on`, eliminating the race condition where routing-fix could start before the file existed.

**Memory cost:** The entire 16,721-entry ipset costs only ~1.6 MiB of kernel memory  negligible in the 600MB VM budget.

#### 15.12.18 Incident Post-Mortem: Discarded Docker Exec Errors in logs (July 10, 2026)
**Symptom:** When the `warp` container is stopped, dead, or failing to connect to its endpoints, the `toggle.sh` and `recover.sh` connectivity health checks fail, but no details are printed to `output.log`. The logs only show generic `FAIL` outputs, making it hard to see if the failure was a network timeout, Docker daemon error, or a missing binary.

**Root Cause:** Overly broad error silencing. The `docker compose exec` commands in the health checks (Check 3 in `toggle.sh` and the traffic check in `recover.sh`) both redirected stderr to `/dev/null` (`&>/dev/null` and `2>/dev/null` respectively), completely throwing away any diagnostic error messages produced by Docker or `wget`.

**Fix (July 10, 2026):** Redirected stderr of the `docker compose exec` check commands to `output.log` (`2>> output.log`). This ensures that any command execution failures or connection timeout details are logged.

#### 15.12.19 Incident Post-Mortem: Toggle Decided STOP/START From a Live Probe, Aborting When Docker Was Down (July 13, 2026)
**Symptom:** Running `./toggle.sh` (or `--restart`) while the Colima VM / Docker daemon was **not** running caused the script to hang for ~15 seconds and then abort during the START phase without ever booting Colima. With `DEBUG_TRACE` the lifecycle breadcrumb showed `toggle.sh EXIT: code=1 phase=start`, and the trace showed execution jumping straight from the START-branch entry to the `ERR` trap (`cleanup_handler ERR`) with the START body never executing. Historically this silent failure occurred 57 in `output.log`. Spam-clicking the toggle during the resulting window (each retry rejected by the lock guard, then finally succeeding) is what produced the earlier "pile of stacked toggle processes" and Wi-Fi-bounce confusion.

**Root Cause:** `toggle.sh` decided whether to STOP or START by calling `is_gateway_active()` (`scripts/common.sh`), which **live-probes** the running system  its first step is `run_with_timeout 15 docker compose ps --status running`. When `/var/run/docker.sock` is absent (Colima down), that call blocks for the full 15 s timeout and returns non-zero. Combined with the script's global `set -e` and `ERR` trap, a non-zero result evaluated at the `if is_gateway_active; then  else  fi` boundary aborted the whole script *before* the `else` (START) branch body ran  precisely the branch whose job is to start Colima. Design-wise the deeper flaw is that a **disable** should not depend on whether the gateway is *currently* healthy; it should act on whether the gateway is *supposed* to be running (persisted intent). The correct signal already existed `MARKER_FILE=/tmp/nullexit-gateway-active.marker`, written at the end of a successful START and cleared on STOP  but `toggle.sh` only wrote/cleared it and never read it for this decision. **Note:** this was a long-standing latent bug, not a regression from the July 2026 `common.sh` platform-function unification or the `DEBUG_TRACE` work; those changes only made the previously-silent abort *visible* (via the `EXEC:`/`EXIT:` breadcrumbs and xtrace).

**Fix (July 13, 2026):** Introduced `gateway_should_be_running()` in `common.sh` and routed all four decision sites through it (`toggle.sh` main decision + `--restart` pre-check; `scripts/toggle-linux.sh` main decision + `--restart` pre-check). It decides from persisted intent:
1. `toggle.sh` captures the marker's existence into `GATEWAY_MARKER_AT_STARTUP` **before** the top-of-script "defensive stale-marker clear" wipes it, and the helper honors that captured value (falling back to the live marker file for callers that run before the clear, e.g. the `--restart` pre-check).
2. Marker present  STOP path; marker absent  START path. This is instant and cannot hang or abort.
3. Only when the marker is absent **and** the Docker socket is actually reachable (`/var/run/docker.sock` or `~/.colima/default/docker.sock` on macOS) does it consult `is_gateway_active` as a non-fatal reconciliation fallback (covers a marker lost to a `/tmp` wipe while containers are genuinely up). A dead socket short-circuits to "stopped" with no probe, so `set -e` can never be tripped.
```

4. The STOP path's `docker compose down` is now `|| true`-guarded so a disable always completes even with Docker down; the START path's `docker compose up` is deliberately left unguarded (a failed start *should* trigger cleanup). On Linux the marker isn't maintained, so the helper degrades to the `is_gateway_active` fallback fast there because native Docker has no Colima VM to hang.

How we got here (methodology note). This one is worth recording because the *process* mattered as much as the patch. The failure had been happening silently for a long time the log simply ended after a teardown with no error, so there was nothing to chase. We first attacked the **observability gap** rather than guessing at the cause: we wrapped the heavyweight lifecycle commands (`colima start`, `docker compose up/down`) so their output and exit codes always land in `output.log`, added an `EXIT: code= phase=` breadcrumb that fires even when the process is killed, and put a full opt-in xtrace behind `DEBUG_TRACE`. Only *then* did we reproduce the failure and this time the trace said, unambiguously, `EXIT: code=1 phase=start` with the START body never entered. The instrumentation didn't fix anything, but it converted an invisible abort into a precise, reproducible fact. (Lesson, consistent with 15.12.18 and the LUKS post-mortem referenced in 15.11: **a mechanism that fails silently is worse than one that fails loudly** so we make it fail loudly first.)

The root-cause leap, though, came from a **design-smell intuition, not the stack trace**: the observation that it is nonsensical for a *disable* to first check whether the gateway is *currently online* a teardown should act on whether the gateway is *supposed* to be running, not probe a subsystem to find out. That instinct turned out to name the bug exactly. The smell was a **conflated responsibility with an inverted dependency**: to decide whether to *start* the gateway, the code first had to *talk to* the gateway (via Docker) the very thing START is responsible for bringing up so the check was guaranteed to be unavailable in precisely the cold-start case that mattered. A second tell was **false symmetry**: enable (builds state) and disable (tears it down) are not symmetric operations, yet both ran the identical expensive probe; whenever two operations that should differ share suspiciously identical machinery, the code is usually lying about the structure of the problem, and reality eventually calls the bluff. The fix invented nothing the honest signal (`MARKER_FILE`, literally "the gateway is supposed to be up") already existed and was maintained in all the right places; it was simply never consulted for the decision it was made for. We just pointed the decision at the answer the codebase already knew. General principle worth keeping: **design smells here are load-bearing they tend to *be* latent bugs, not merely cosmetic**, and are worth chasing on intuition even before a trace confirms them.

Acceptance verified (static + live): with marker present STOP (instant, even with Docker down); with marker absent + Docker down START (0 s, no `set -e` abort); a real run now proceeds past init to the gateway-status decision. `bash -n` clean on all three scripts; no new `shellcheck` findings; re-signed `.signatures` (`toggle.sh`/`common.sh`/`toggle-linux.sh` are in the signed set) and `crypto.sh --verify` exits 0.

Follow-up (same day) a `set -e` regression the fix introduced, and why only a live run caught it. The first cut of the marker capture wrote the intent as a full `if [ -f "$MARKER_FILE" ]; then GATEWAY_MARKER_AT_STARTUP="yes"; else GATEWAY_MARKER_AT_STARTUP="no"; fi`. On macOS `/bin/bash` 3.2, with `set -e` active and the `DEBUG_TRACE` xtrace redirect (`exec 2> >(tee )`) in effect, that construct **aborted the script at init** the moment the marker was absent the false `[ -f ]` status leaked through the compound and `set -e` killed the run (`EXIT: code=1 phase=init`), before the gateway logic ran at all. Notably this did **not** reproduce in isolated `bash -c` snippets of the same block it only surfaced in the script's real runtime context so we had to bisect against the actual `toggle.sh` (swap the block for a trivial assignment the script sailed past init) to prove causation. Fix: write it as a default assignment plus an `if` **without an `else`** (`GATEWAY_MARKER_AT_STARTUP="no"` then `if [ -f "$MARKER_FILE" ]; then GATEWAY_MARKER_AT_STARTUP="yes"; fi`), which yields status 0 when the condition is false. Lesson reinforced: `set -e` on legacy bash is genuinely treacherous around conditionals, and the observability work paid for itself again the breadcrumb pinned the abort to init instantly, and a live run (not static checks) was the only thing that exposed a heisenbug invisible to isolated reproduction.

Part IV Observations, Changelog & TODO

16. Unverified Observations

16.1 AdGuard Returns Two Different Sinkhole IPs (Unverified)

> **Status: Observed but not formally verified. Needs a deeper audit of the rule-compiler sources to confirm.**

Observation (July 10, 2026): When querying AdGuard Home for blocked ad domains, some domains resolve to `0.0.0.0` and others resolve to `127.0.0.1`. Both are blocked, both cause connection failure, but the different responses were unexpected.

```
doubleclick.net              0.0.0.0    (blocked)
ads.google.com               127.0.0.1  (blocked)
pagead2.googlesyndication.com  0.0.0.0  (blocked)
scorecardresearch.com        127.0.0.1  (blocked)
```

Likely Explanation (Unverified): There are three historical schools of thought on the "correct" DNS sinkhole response, and AdGuard faithfully preserves whichever the source blocklist used:

Sinkhole IP	Convention	Origin

```
| `0.0.0.0` | AdBlock-style lists (`\|\|domain^`) | Modern DNS-level blocking; AdGuard's native format. Fails fast  OS doesn't attempt a TCP connection. Works for both IPv4 and IPv6 without a separate rule. |
| `127.0.0.1` | Hosts-file-style lists (`127.0.0.1 domain`) | 90s-era `/etc/hosts` tradition. Steven Black's hosts list (500k+ domains) still uses this. |
| `NXDOMAIN` | Purist / Pi-hole default | Returns "domain doesn't exist." Semantically most accurate but requires browsers to handle NXDOMAIN gracefully. |

The dual-response behavior in this project is because `rule-compiler` pulls from both AdBlock-format sources (Hagezi, uBlock origin lists) and hosts-format sources (Spamhaus DROP, Steven Black, Feodo Tracker), and AdGuard returns whatever sinkhole IP the matching rule specifies.

**To Verify:** Run `docker compose exec tailscale cat /etc/hosts` to confirm no hosts-file rules are being injected at the container level, and check which specific AdGuard rule matched each domain via the AdGuard query log at `http://100.100.21.8:3000/#logs`.

### 16.2 AGY CLI Appears to Generate Native macOS Notification Pop-ups (Unverified)

> **Status: Observed but source unverified. macOS security logs ruled out system-level origin.**

**Observation (July 10, 2026):** While running a DNS blocklist verification via the Antigravity CLI (`agy`), a command was proposed that included `malware.wicar.org` (a known-safe DNS blocklist test domain) in a `dig` loop. Shortly after, a macOS-style notification popup appeared describing a "malicious" or "blocked" event. The AGY CLI terminal session then closed.

**Investigation:**
- **macOS XProtect / Gatekeeper logs:** Empty. Zero events in the 30-minute window around the incident.
- **macOS Endpoint Security / System Policy logs:** Empty.
- **Broad `log show` search for "malicious", "malware", "blocked":** Empty.
- **Conclusion:** macOS itself did not generate the notification. No system security subsystem fired.

**Likely Explanation (Unverified):** The notification appears to have originated from **AGY CLI's own internal safety system**. AGY CLI is a native macOS application (not just a terminal process) and has access to the macOS Notification Center API (`UNUserNotificationCenter`). When its guardrails detected a domain containing the word "malware" (`malware.wicar.org`) in a proposed shell command, it likely:
1. Blocked the command from executing
2. Posted a native macOS-style notification via the Notification Center

This is surprising behavior for a CLI tool  most terminal programs don't send GUI notifications  but AGY CLI appears to bridge both worlds: it runs in a terminal but is packaged as a native app with full system API access. The notification would be indistinguishable from a system security alert at a glance.

**Notes:**
- `malware.wicar.org` is the DNS/AV equivalent of the EICAR test file  a deliberately benign domain that triggers blocklist/AV systems to verify they work. It was included in the test set to confirm the AdGuard blocklist was catching it. AGY's safety system is not aware of this distinction.
- For future DNS blocklist testing, use only ad/tracking domains (e.g., `doubleclick.net`, `adnxs.com`) to avoid triggering AGY's guardrails.

### 16.3 `system_profiler SPAirPortDataType` Longevity (Wi-Fi Detection)
The Wi-Fi security detection used by `watcher.sh`'s `detect_lan_p2p_mode()` currently relies on `system_profiler SPAirPortDataType` (after the `airport` binary removal  full history in 15.11.6). This works as of macOS Sequoia (15.x) without sudo and without redaction, but given Apple's ongoing Wi-Fi privacy clampdown, it may be locked down in a future release. If `security` comes back empty on a future macOS version while on Wi-Fi, detection silently falls back to `allow=false` (a safe default). The long-term fix would be a small Swift CoreWLAN helper compiled and bundled in the repo.

## 17. Changelog / Recent Updates

Reflects the current branch's state. Each entry is a one-line summary + commit hash; read the commit message (and the linked Incident Log entry) for full detail. Full post-mortems live in 15.

- **July 10, 2026  `fix(race): suppress WARP Watcher during post-wake force-recreate`**  Fixed a race where switching Wi-Fi caused the host IP to leak as the raw ISP IP (`132.x`) on every roam; `recover.sh --post-wake`'s warp force-recreate tripped the WARP Watcher into a nuclear shutdown. Fixed via `/tmp/nullexit-warp-inhibit.marker`. See 15.2.7.
- **July 10, 2026  startup/roaming stabilization suite**  Pre-flight Tailscale race (15.5.3), STUN-block cold-boot false negative (15.5.4), restart DNS-build race via `wait_for_dhcp_settle` (15.5.5), nuclear/post-wake concurrency lock (15.2.8), infinite auto-recovery loop marker leak (15.2.9), PF kill-switch bricking SOCKS5 fallback (15.7.3), `--reset` host logouts (15.7.4), missing-flags abort (15.7.5), discarded docker-exec errors (15.12.18), and the SNAT endpoint poisoning fix (15.7.1).
- **July 6, 2026  `fix: resolve edge-case vulnerabilities and system state leaks`**  Security-review hardening suite:
  1. **DNS Watcher Interface Independence:** `start_dns_watcher` now detects the active interface via `get_active_service()` instead of hardcoding a `grep` for "Wi-Fi".
  2. **Robust Lock Verification:** `toggle.sh` lock-file logic uses `ps -p` verification to prevent permanent lockouts from a recycled PID.
  3. **Fail-Closed Crypto:** `crypto.sh` integrity check switched from `[ -x ]` to `[ -f ]` so it runs even if git strips the `+x` bit.
  4. **Strict `iptables` Pinning:** `post-rules.txt` FORWARD-chain rules explicitly pin `tailscale0``tun0` both directions, preventing escape via `eth0` during WARP drops.
```

```
  5. **Docker Local Network Isolation:** `docker-compose.yml` binds exposed WARP ports (`1080`, `5354`) to `127.0.0.1` to prevent LAN access.
  6. **Routing Verification & Sudoers:** Added a `netstat -nr` default-route override check and ensured `/usr/bin/python3` is permitted in the sudoers template.
- **July 4, 2026  `cc789c5`**  Concurrency mutex on `toggle.sh` (`/tmp/nullexit-toggle.lock`); defensive `TS_AUTHKEY` init under `set -u`; Go `go vet`/`go build` steps in CI; AdGuard log capping; macOS BSD `grep` fixes in `setup-common.sh`; `recover-linux.sh` quoting/PID fixes; removed redundant bypass-routes block in `recover.sh`.
- **July 4, 2026  `fix: post-wake routing loop & destructive recovery`**  (1) `add_warp_bypass_routes` now accepts explicit interface args so post-wake stops routing WARP traffic back into `utun*` (was an infinite loop killing internet on every wake). (2) Replaced `sudo route -n flush` with targeted deletions of `0.0.0.0/1`/`128.0.0.0/1` + WARP bypass routes, so recovery no longer destroys loopback/multicast routes and wedges `configd`/`mDNSResponder` until reboot.
- **July 4, 2026  `feat: secure execution`**  Restricted the blanket `/usr/bin/python3` NOPASSWD sudo to exactly `/usr/bin/python3 -I .../scripts/dns-proxy.py`; locked `dns-proxy.py` with `chown root:wheel` + `chmod 755`; added native auth prompts to `Toggle Gateway.app`/`Recover Gateway.app`; `toggle.sh` now captures `warp` logs on failure before teardown.
- **July 4, 2026  Go ports + refactor**  `logger` and `rule-compiler` ported to Go multi-stage Docker builds (15.8.6); ~200 lines of duplicated bash extracted to `scripts/common.sh` (15.9.4); `crypto.sh` HMAC-SHA256 script integrity added.
- **July 3, 2026  `fix: resolve numerous system regressions`**  Fixed truncated `/etc/sudoers.d/nullexit` line; temporary `.tmp`+`os.replace()` atomic writes in `sync-rules.py` (later reverted, see below); `recover.sh` `ts_args` quoting; ANSI `%b`/`-e` output in `diagnose-host-leak.sh`/`fix-docker-bridge-collision.sh`; `toggle-linux.sh` using `resolvectl dns` instead of macOS `networksetup`; AdGuard REST refresh POST targeting port `3000`; restored `START_GATEWAY=true`; purged stale port `80` mapping; verified `output.log` in `.gitignore`.
- **July 3, 2026  `unlock-files.sh` + revert atomic writes**  Reverted all 5 `.tmp`+`os.replace()` sites in `sync-rules.py` to direct writes; created `scripts/unlock-files.sh` to permanently fix stale `0444`/`000` file permissions via atomic rename (no `chmod`). Covers `.env`, `compiled_rules.txt`, `ip_blocklist.ipset`, `cache/*.txt`, `cache/ip/*.txt`, `data/filters/*.txt`. See 15.9.6 + 3.
- **July 3, 2026  Overnight leak + geo-blocking**  Introduced `scripts/host-leak-probe.sh` (sub-second host-egress detector, 15.6.1) and the statistical-threshold WARP auto-shutdown watcher (15.6.2); fixed the overnight silent IP leak (15.6.3).
- **July 2, 2026  `8c5fae1` + `181e1e1`**  Two infinite routing loops fixed: host-VM tunnel loop (WARP bypass routes via physical gateway IP + `setup_exit_node_routing` forcing default to `utun*`) and Docker Compose subnet takeover (`10.200.1.0/24` lock). `--accept-routes=true` added explicitly to all `tailscale up --reset` calls. See 15.2.2.
- **July 2, 2026  auto-recovery daemon**  `scripts/watcher.sh` + launchd LaunchAgent documented. See 15.5.1.
- **July 1, 2026  `c5b92c0` (refactor: personal-leak cleanup)**  Eliminated every residual personal-username leak across source + config; launchd plist uses `WorkingDirectory` + relative program arg with the `__NULLEXIT_HOME__` sentinel (install recipe gains a single `sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|"` step); 8 devref prose sites genericized to `~/.colima/...`, `~/Library/LaunchAgents/...`, `$USER`, `<USER>`.
- **July 1, 2026  `a5e2a5a` (refactor: script-relative paths)**  Replaced 6 hardcoded `/Users/<user>/.../nullexit/...` literals with script-relative resolution. `recover.sh`/`watcher.sh` inject `SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"`; `RECOVER="$SCRIPT_DIR/../recover.sh"` resolves without launch-path dependency.
- **June 28, 2026  `f8f8e4e` (M1: scripts/ move)**  8 internal scripts consolidated into `scripts/`; the 4 user-facing/orchestrator/config files (`toggle.sh`, `recover.sh`, `setup.sh`, `docker-compose.yml`) stayed at repo root.
- **June 28, 2026  `0875ef3` (ci: glob)**  CI `py_compile` switched to `python3 -m py_compile scripts/*.py` after M1 left root with zero `.py` files.
- **June 28, 2026  `25b37a1` (docs: scripts/ prefix)**  `devref.md` + `README.md` updated to the `scripts/` prefix for all 8 moved files (~31 patches).

### End-to-end verification (commit `c5b92c0`)
Manual START  STOP cycle ran clean on macOS after path relativization + personal-leak cleanup:
- `bash toggle.sh` START: exit 0, ~135s (cold Colima boot); marker present (`2026-07-02T03:41:41Z`), all 5 containers healthy, host DNS hijacked to `100.100.21.8` (no fallback), exit node selected, Cloudflare trace through WARP succeeds.
- `bash toggle.sh` STOP: exit 0, ~12s; marker cleared, all nullexit containers gone, host DNS restored to `1.1.1.1`, Tailscale disconnected.

## 18. TODO & Future Work

### 18.1 TODO
* **Decide toggle STOP/START from persisted *intent*, not a live probe:** ~~`toggle.sh` decided STOP-vs-START by calling `is_gateway_active()`, which live-probes the running system (`docker compose ps` + DNS + SOCKS checks). When Colima/Docker was down the `docker compose ps` blocked for the full 15 s `run_with_timeout` window and, under `set -e` + the `ERR` trap, aborted the script *before the START body ran*  so the path that would boot Colima never executed (`EXIT: code=1 phase=start`, seen 57 historically).~~ **Closed** added `gateway_should_be_running()` in `common.sh`, which reads persisted intent (`MARKER_FILE`) captured into `GATEWAY_MARKER_AT_STARTUP` before the defensive stale-marker clear, and only falls back to `is_gateway_active` when the marker is absent *and* the Docker socket is actually reachable (so a dead-socket probe can neither hang nor trip `set -e`). All four decision sites (`toggle.sh` main + `--restart`, `toggle-linux.sh` main + `--restart`) now use it, and the STOP `docker compose down` is `|| true`-guarded so a disable always completes even with Docker down. See 15.12 Misc Resolved.
* **Verify Wi-Fi Roaming:** Investigate and test whether switching Wi-Fi networks now successfully and smoothly recovers the gateway end-to-end without any tailscale flag errors or dead tunnels. (Pending user verification).
* **Fix Diagnostic Script Check 5/8:** ~~Investigate why `diagnose-host-leak.sh` check 5/8 (`Host default route`) sometimes reports "default route goes via physical Wi-Fi  Tailscale route assertion failed" on macOS
```

```
     even though the `warp=on` egress checks confirm that traffic is successfully tunneling.~~ **Closed**  fixed
      by switching check 5 from `netstat -rn | grep default` to `route -n get 1.1.1.1`. macOS Tailscale uses
     interface-scoped routes and does not replace the DHCP default route entry. See 15.7.1 Known Unknowns.
* **Expose Tor Control Port (9051):** ~~Consider mapping the Tor Control port to localhost and adding `nyx` (
     the Tor terminal status monitor) to allow users to inspect their entry, middle, and exit node IPs in real-
     time.~~ **Closed**  mapped the Tor Control Port to a procedurally generated, localhost-bound port `
     TOR_CONTROL_PORT`, configured `ControlPort` and disabled cookie authentication in `docker/tor/entrypoint.sh
     `, and integrated it into the `/sweep` verification protocol.
* **Host-Only Tor Mode (Pluggable Egress):** Implement an `EGRESS_TYPE=tor_host_only` mode for users in heavily
      censored (DPI-restricted) environments. This mode would skip booting `warp` and `tailscale`, boot `
     adguardhome` and `tor` (using Telegram `obfs4` bridges), set AdGuard's upstream to Tor's `DNSPort`, and use
      the `pf` kill-switch to force all host Mac traffic strictly through the Tor network. This provides a 100%
     free, DPI-resistant evasion path without requiring an external VPS.
  * **The Technical Realities (What You Need to Watch Out For):**
    * **The UDP Problem:** Tor only transports TCP traffic. It does not support UDP (aside from DNS requests
     passed directly to its DNSPort). macOS is incredibly chatty over UDP (FaceTime, mDNS, QUIC protocols). Your
      pf rules defined in `scripts/pf.conf` must be configured to aggressively and silently DROP all host UDP
     traffic (except for local DNS queries to AdGuard). If you don't drop it cleanly, macOS services might hang
     indefinitely while waiting for timeouts.
    * **The "Daily Driver" Friction:** Routing a whole host OS through Tor is brutal on performance. Background
      daemons (iCloud sync, macOS software updates, Spotlight indexing) will blindly attempt to download
     gigabytes of data over Tor, which could saturate the circuits or time out completely.
    * **Bridge Rot:** State firewalls actively hunt and block obfs4 bridges. When a user's bridges burn out,
     they will lose all internet access due to the pf kill-switch. You will need to ensure the user has a
     frictionless way to update the `tor-bridges.txt` file and restart the Tor container without exposing their
     IP.
    * **Exit Node Discrimination:** Because all traffic will exit from known Tor nodes, the user will face an
     onslaught of Cloudflare CAPTCHAs, and many banking or streaming services will flat-out reject the
     connections.

### 18.2 Future Work
- **Direct P2P on VM Hosts:** Non-Linux hosts run Docker inside a VM, making true UDP hole-punching **to
     arbitrary external peers** impossible; those peers fall back to the DERP relay unless bridged mode is
     enabled on a trusted LAN. The only fully general solution is a native Linux host (Raspberry Pi, Intel NUC)
     where Docker runs without VM hypervisor translation. *(The one case that works regardless is direct **
     hostcontainer** P2P  the gateway and its own Mac, the same physical machine  re-permitted via routing-fix
     `.host_ips` RETURN rules and confirmed `direct` even in plain user-mode networking; see 15.3.5 and 15.8.7.)
     *
- **Native Linux Deployment:** Benchmark on a Raspberry Pi  native container footprint is ~75MB without macOS
     hypervisor overhead.
- **Mesh-Wide Filesystem (SFTP over Tailscale):** Any mesh device passively exposes its filesystem over SFTP.
     Android via Termux + OpenSSH + wakelock; iOS is client-only due to background process limits.
- **Post-Quantum Cryptography (PQC):** Tailscale uses Curve25519, theoretically vulnerable to "harvest now,
     decrypt later" quantum attacks. A future iteration could replace the Tailscale container with raw WireGuard
      + [Rosenpass](https://rosenpass.eu/) to negotiate post-quantum PSKs.
- **Egress Compartmentalization (Pluggable Architecture):** See 12.9 for the full plan to make the entire
     egress layer swappable via a single `.env` variable.
```

5 diagrams.md

```
# nullexit  Flow Diagrams & Architecture

---

## 1. System Architecture  What's Running & Where

```mermaid
graph TB
    subgraph INTERNET["  Internet / Cloudflare Edge"]
        CF["Cloudflare WARP\n(Exit Point)\nwarp=on"]
    end

    subgraph TAILNET[" Tailscale Mesh (100.x.x.x)"]
        PHONE[" Phone / Laptop\n(Tailnet Client)"]
    end

    subgraph MACHOST[" macOS Host"]
        direction TB
        PFFIREWALL["macOS pf Firewall\n(Kill-Switch & TCP MSS Clamping)"]
        TSDAEMON["tailscaled\n(brew service)\nhost Tailscale"]
        HOSTDNS["Host DNS\n Gateway IP\n(hijacked by toggle.sh)"]
        HOSTSOCKS["Host SOCKS5\n127.0.0.1:1080\n(fallback only)"]
        HOSTTOR["Host Tor SOCKS5\n127.0.0.1:$TOR_SOCKS_PORT\n(randomized)"]
        HOSTTORCTRL["Host Tor Control\n127.0.0.1:$TOR_CONTROL_PORT\n(randomized)"]

        subgraph MONITORS[" Background Daemons (toggle.sh children)"]
```
```

```
            WARPWATCH["WARP Watcher\n(every 30s, via docker exec)\n output.log"]
            LEAKPROBE["Host Leak Probe\n(every 300ms, via curl from HOST)\n output.log"]
            DNSWATCHER["DNS Watcher\n(re-hijacks if drift detected)\n output.log"]
            CAFFEINATE["caffeinate -i\n(sleep prevention)"]
        end

        subgraph COLIMAVM["  Colima VM (Linux, 600MB RAM)"]
            subgraph NETNS["warp container network namespace (shared by ALL containers)"]
                GLUETUN["warp / Gluetun\nWireGuard  tun0\n(owns the namespace)"]
                TS["tailscale container\nAdvertises exit node\ntailscale0 interface"]
                SOCKS["tailscale\nSOCKS5\nport 1080"]
                ADGUARD["adguardhome\nDNS sinkhole\nport 5335"]
                ROUTINGFIX["routing-fix sidecar\nGeo-IP Blocking & Go Logger\n(writes blocked.log)"]
                TOR["tor container\nSOCKS5: port 9050\nControl: port 9051\nTransparent Proxy: port 9040\
    nDNSPort: port 5353"]
            end
        end

        subgraph LAUNCHD[" launchd (always-on)"]
            WATCHER["scripts/watcher.sh\n(sleep/wake + Wi-Fi roam\n+ LAN P2P auto-detection)"]
        end
    end

    subgraph RECOVERY[" Recovery"]
        RECOVER["recover.sh\n(nuclear reset)"]
        RECOVERPOST["recover.sh --post-wake\n(light refresh)"]
    end

    %% Traffic flow
    PHONE -->|"WireGuard / Tailscale mesh"| TSDAEMON
    TSDAEMON -->|"exit-node route utun*"| TS
    TS -->|"FORWARD  tun0 (MASQUERADE)"| GLUETUN
    GLUETUN -->|"WireGuard UDP (macOS Host Egress)"| PFFIREWALL
    PFFIREWALL -->|"TCP MSS Clamped / Kill-Switch check"| CF
    TS -->|"198.18.0.0/15 PREROUTING redirect"| TOR
    TOR -->|"internal path / exit nodes via tun0"| GLUETUN

    %% DNS
    HOSTDNS -->|"UDP:53  TCP:5354"| ADGUARD
    ADGUARD -->|"DNS upstream via tun0"| CF
    ADGUARD -->|"onion queries  localhost:5353"| TOR

    %% Monitoring
    WARPWATCH -->|"docker exec warp wget cdn-cgi/trace"| GLUETUN
    LEAKPROBE -->|"curl cdn-cgi/trace (HOST NIC)"| CF
    DNSWATCHER -->|"networksetup -getdnsservers"| HOSTDNS

    %% Logging
    ROUTINGFIX -->|"writes"| BLOCKEDLOG["blocked.log\n(Host-side)"]
    HOSTTOR -->|"port mapping"| TOR
    HOSTTORCTRL -->|"port mapping"| TOR

    %% Recovery triggers
    WARPWATCH -->|"6 consecutive warp=off"| RECOVER
    WATCHER -->|"sleep/wake / roam"| RECOVERPOST
    RECOVERPOST -->|"if gateway unhealthy"| RECOVER
    WATCHER -->|"writes .lan_p2p_detected"| ROUTINGFIX

    style MONITORS fill:#1a1a2e,color:#e0e0ff
    style COLIMAVM fill:#0d2137,color:#e0e0ff
    style NETNS fill:#0a1628,color:#e0e0ff
    style RECOVERY fill:#2d0a0a,color:#ffe0e0
    style INTERNET fill:#0a2d1a,color:#e0ffe0
    style TAILNET fill:#1a2d0a,color:#e8ffe0
    style LAUNCHD fill:#2d1a2d,color:#ffe0ff
```

---

2. toggle.sh Full START Flow

```mermaid
flowchart TD
    START(["./toggle.sh"])
    CHECK{"is_gateway_active?\n(containers running OR\nDNS  1.1.1.1)"}

    START --> CHECK
    CHECK -->|"YES  already ON"| STOPFLOW[" See STOP Flow"]
    CHECK -->|"NO  start it"| S1

    S1["1. Reset DNS  1.1.1.1\n(prevent deadlocks)"]
```

```
    S2["2. tailscale down\n(prevent exit-node routing during boot)"]
    S3["3. Check Colima VM\n(colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged)"]
    S4["4. Configure VM swap\n(400MB, prevent OOM)"]
    S5["5. Clean corrupted AdGuard config\n(prevent container crash loop)"]
    S6["6. docker compose up -d --build\n(compile Go threat rules & boot containers)"]
    S7["7. Poll tailscale status\n(wait for 'offers exit node'\nup to 60s)"]
    TSREADY{"container\nTailscale ready?"}
    ABORT([" ABORT\n cleanup_handler"])
    S8["8. Resolve gateway Tailscale IP\n(.gateway_ip or docker exec)"]
    S9["9. Verify tailscaled running\n(auto-start if needed)"]
    S10["10. tailscale up --reset\n--exit-node=\n--accept-routes=true\n--ssh=true --accept-dns=false"]

    PF1{"Pre-flight 1/3\ntailscale ping gateway"}
    PF2{"Pre-flight 2/3\ndig +tcp AdGuard DNS"}
    PF3{"Pre-flight 3/3\nWARP internet check"}
    ALLPASS{"All 3 pass?"}

    EXITPATH[" EXIT NODE PATH\n tailscale up --exit-node=\n force_dns_to_gateway (3 retry)\n SOCKS5 disabled"]
    SOCKSPATH[" SOCKS5 FALLBACK PATH\n macOS SOCKS5  127.0.0.1:1080\n DNS proxy  127.0.0.1:53\n No exit node (
     SKIP_EXIT_NODE=true)"]

    DAEMONS["Start background daemons\n caffeinate -i (sleep prevention)\n DNS Watcher\n WARP Watcher\n Host
     Leak Probe\n(all write  output.log)"]

    RULECOMPILE["Rule compiler status\n(log rule/IP counts)"]
    WRITEMARKER["write_gateway_active_marker\n(/tmp/nullexit-gateway-active.marker)"]
    DONE([" Gateway UP"])

    S1 --> S2 --> S3 --> S4 --> S5 --> S6 --> S7
    S7 --> TSREADY
    TSREADY -->|"NO (40 NoState)"| ABORT
    TSREADY -->|"YES"| S8
    S8 --> S9 --> S10
    S10 --> PF1 --> PF2 --> PF3
    PF1 & PF2 & PF3 --> ALLPASS
    ALLPASS -->|"YES"| EXITPATH
    ALLPASS -->|"NO"| SOCKSPATH
    EXITPATH --> DAEMONS
    SOCKSPATH --> DAEMONS
    DAEMONS --> RULECOMPILE --> WRITEMARKER --> DONE

    style ABORT fill:#5c1010,color:#fff
    style EXITPATH fill:#0a3d1a,color:#c0ffc0
    style SOCKSPATH fill:#3d2d00,color:#ffe0a0
    style DAEMONS fill:#0a1f3d,color:#c0d8ff
    style DONE fill:#0a3d1a,color:#fff
```

---

3. toggle.sh STOP Flow

```mermaid
flowchart TD
    STOP(["./toggle.sh\n(gateway already ON)"])

    ST1["1. tailscale down\n(disconnect from mesh)"]
    ST2["2. docker compose down -t 5\n(stop all containers)"]
    ST3["3. Reset DNS  1.1.1.1\n(networksetup on all services)"]
    ST4["4. stop_local_dns_proxy\n(kill Python UDP proxy)"]
    ST5["5. stop_pidfile_daemon\n(kill caffeinate, rm PID)"]
    ST6["6. stop_pidfile_daemon\n(kill DNS Watcher, rm PID)"]
    ST7["7. stop_pidfile_daemon\n(kill WARP Watcher, rm PID)"]
    ST8["8. stop_pidfile_daemon\n(kill Leak Probe, rm PID)"]
    ST9["9. clear_gateway_active_marker\n(rm /tmp/nullexit-gateway-active.marker)"]
    ST10["10. cleanup_network_state\n disable_all_proxies\n Flush DNS cache (dscacheutil)\n Flush stale routes\
     n Power-cycle Wi-Fi (offon)"]

    DONE([" Gateway DOWN\nInternet restored"])

    STOP --> ST1 --> ST2 --> ST3 --> ST4
    ST4 --> ST5 --> ST6 --> ST7 --> ST8 --> ST9 --> ST10 --> DONE

    style DONE fill:#0a3d1a,color:#fff
```

---

4. Monitoring Layer The 4 Background Daemons

```mermaid
```

```
graph LR
    subgraph DAEMONS["Background Daemons (all start with gateway, stop on teardown)"]
        direction TB

        subgraph D1[" Sleep Prevention"]
            CAFF["caffeinate -i\nPID: /tmp/nullexit-caffeinate.pid\n\nKeeps Mac awake while\ngateway is active
    "]
        end

        subgraph D2[" DNS Watcher"]
            DNSW["Polls networksetup every ~30s\nPID: /tmp/nullexit-dns-watcher.pid\n\nIf DNS drifts from
    gateway IP\n re-hijacks automatically\n output.log"]
        end

        subgraph D3[" WARP Watcher"]
            WARPW["Polls every 30s via:\ndocker compose exec warp wget cdn-cgi/trace\nPID: /tmp/nullexit-warp-
    watcher.pid\n\nCounts consecutive warp=off\n WARP_FAIL_THRESHOLD (default 6=30s)\n runs recover.sh (nuclear
    )\n output.log"]
        end

        subgraph D4[" Host Leak Probe"]
            LEAKP["Polls every 300ms via:\ncurl cdn-cgi/trace (HOST NIC directly)\nPID: /tmp/nullexit-host-leak
    -probe.pid\n\nLogs ONLY on state change:\n  LEAK / ROTATE / HOST-PROBE failed\nAll errors  output.log\n(
    HOST_LEAK_PROBE=true in .env)"]
        end
    end

    subgraph BLIND["What each monitor can/can't see"]
        B1["WARP Watcher sees:\n Container WARP tunnel state\n Host NIC traffic\n Flash leaks < 5s"]
        B2["Host Leak Probe sees:\n Host NIC egress (browser path)\n Sub-second transitions\n Container
    internals"]
    end

    D3 -.->|"covers"| B1
    D4 -.->|"covers"| B2

    style D1 fill:#1a2d0a
    style D2 fill:#0a1a2d
    style D3 fill:#2d0a2d
    style D4 fill:#2d1a00
    style BLIND fill:#1a1a1a,color:#aaa
```

---

5. Traffic Flow Data Path (Normal Operation)

```mermaid
sequenceDiagram
    participant PHONE as  Phone<br/>(Tailnet client)
    participant TS_HOST as tailscaled<br/>(macOS host)
    participant TS_CONTAINER as tailscale<br/>(container)
    participant GLUETUN as Gluetun/warp<br/>(container)
    participant CF as  Cloudflare<br/>WARP Edge
    participant INTERNET as  Internet

    Note over PHONE,INTERNET: Normal EXIT NODE path (all 3 pre-flights passed)

    PHONE->>TS_HOST: WireGuard packet<br/>(encrypted, UDP 41641)
    TS_HOST->>TS_CONTAINER: route via utun*  tailscale0
    TS_CONTAINER->>GLUETUN: FORWARD chain<br/>(iptables MASQUERADE on tun0)
    GLUETUN->>TS_HOST: WireGuard tunnel (tun0) egresses macOS Host
    Note over TS_HOST: macOS pf Firewall applies<br/>Kill-Switch & TCP MSS Clamping
    TS_HOST->>CF: re-encrypted UDP packet (MSS 1160)
    CF->>INTERNET: Plain HTTPS from<br/>Cloudflare IP

    Note over PHONE,INTERNET: DNS Resolution path
    PHONE->>TS_HOST: DNS query
    TS_HOST->>TS_CONTAINER: UDP:53  TCP:5354  AdGuard :5335
    TS_CONTAINER->>GLUETUN: AdGuard upstream DNS via tun0
    GLUETUN->>CF: Encrypted DNS query
    CF-->>GLUETUN: DNS response
    GLUETUN-->>TS_CONTAINER: response
    TS_CONTAINER-->>TS_HOST: response
    TS_HOST-->>PHONE: DNS answer
```

---

6. recover.sh Decision Tree

```mermaid
flowchart TD
    INVOKE(["recover.sh called"])
    MODE{"--post-wake\nflag?"}

    subgraph POSTWAKE[" --post-wake (light refresh, gateway stays UP)"]
        PW1["Re-hijack DNS  gateway IP"]
        PW2["Flush DNS cache"]
        PW3["Refresh tailscale exit-node"]
        PW4["Force-recreate warp container\n(if Gluetun healthcheck failed)"]
        PW5["Leave: containers, Wi-Fi,\ncaffeinate, DNS watcher,\nHost Leak Probe  untouched"]
    end

    subgraph NUCLEAR[" Default (nuclear  gateway torn down)"]
        N0["Sudo keep-alive loop (background)"]
        N1["1. tailscale down\n(disconnect from mesh)"]
        N2["2. Reset DNS  1.1.1.1\n(ALL network services)"]
        N3["3. disable_all_proxies\n(all interfaces)"]
        N4["4. Flush DNS cache\n(dscacheutil -flushcache)"]
        N5["5. Flush stale routes\n(WARP endpoint routes)"]
        N6a["6a. docker compose down -t 30\n(stop all containers)"]
        N6b["6b. stop_pidfile_daemon\n(kill caffeinate)"]
        N6c["6c. stop_pidfile_daemon\n(kill DNS Watcher)"]
        N6d["6d. stop_pidfile_daemon\n(kill Leak Probe)"]
        N7["7. Power-cycle Wi-Fi\n(networksetup offsleep 2on)"]
        N8["8. Verify internet working\n(curl ifconfig.io)"]
    end

    DONE_PW([" Gateway refreshed\n(still running)"])
    DONE_NUC([" Internet restored\n(gateway fully stopped)"])

    INVOKE --> MODE
    MODE -->|"yes"| PW1
    PW1 --> PW2 --> PW3 --> PW4 --> PW5 --> DONE_PW
    MODE -->|"no"| N0
    N0 --> N1 --> N2 --> N3 --> N4 --> N5
    N5 --> N6a --> N6b --> N6c --> N6d --> N7 --> N8 --> DONE_NUC

    style POSTWAKE fill:#0a2d1a,color:#c0ffc0
    style NUCLEAR fill:#2d0505,color:#ffc0c0
    style DONE_PW fill:#0a3d1a,color:#fff
    style DONE_NUC fill:#3d1a00,color:#ffe0c0
```

---

7. Failure Paths & Self-Healing

```mermaid
flowchart LR
    subgraph EVENTS["Failure / Event"]
        E1["WARP tunnel drops\n(warp=off)"]
        E2["Mac wakes from sleep\nor Wi-Fi roams"]
        E3["DNS drifts from\ngateway IP"]
        E4["toggle.sh crashes /\nCtrl-C / SIGTERM"]
        E5["Host-side flash leak\ndetected (HOST NIC)"]
        E6["Silent NAT timeout\n(ping stops working)"]
    end

    subgraph DETECTORS["Detector"]
        D1["WARP Watcher\n(30s poll, docker exec)"]
        D2["scripts/watcher.sh\n(launchd, always on)"]
        D3["DNS Watcher\n(30s poll, networksetup)"]
        D4["cleanup_handler\n(ERR/INT/TERM/HUP trap)"]
        D5["Host Leak Probe\n(300ms poll, HOST curl)"]
        D6["routing-fix.sh\n(30s curl over tun0)"]
    end

    subgraph ACTIONS["Response"]
        A1[" threshold consecutive off\n recover.sh (nuclear)"]
        A2["recover.sh --post-wake\n(gentle refresh)"]
        A3["Re-hijack DNS\n(networksetup setdnsservers)"]
        A4["Stop all daemons\nReset DNS  1.1.1.1\nTear down Tailscale"]
        A5["Log LEAK to output.log"]
        A6["Blackhole internet traffic\n+ Drop TUNNEL_FAILED_CLOSED.marker"]
        A7["Push macOS Desktop Notification\n(via watcher.sh listener)"]
    end

    E1 --> D1 --> A1
    E2 --> D2 --> A2
    E3 --> D3 --> A3
```

```
    E4 --> D4 --> A4
    E5 --> D5 --> A5
    E6 --> D6 --> A6
    A6 -.->|"marker detected"| D2
    D2 --> A7

    style EVENTS fill:#2d1010,color:#ffcccc
    style DETECTORS fill:#0a1a2d,color:#cce0ff
    style ACTIONS fill:#0a2d0a,color:#ccffcc
```

---

8. output.log Who Writes What

All events converge into a single `output.log` at the repo root.

Writer	Event Types	Format
`toggle.sh`	All startup/shutdown steps, errors, pre-flight results	Plain text with timestamps
`recover.sh`	Every recovery step (nuclear or post-wake)	Plain text
WARP Watcher	`WARP DOWN`, `WARP RECOVERED`, `WARP SHUTDOWN`	`[UTC timestamp]` prefix
Host Leak Probe	`LEAK`, `ROTATE`, `HOST-PROBE failed/timeout`	`[HH:MM:SS.mmm]` prefix
DNS Watcher	DNS re-hijack events	Appended inline
`docker compose`	Container logs on failure (last 100 lines of warp)	Dumped on ERR path
`routing-fix.sh`	(via `nullexit-logger` inside container)	Structured

Grep cheatsheet:
```bash
grep LEAK output.log           # Host-side warp=off events (real leaks)
grep ROTATE output.log         # Cloudflare edge IP rotations (harmless)
grep 'HOST-PROBE' output.log   # curl failures from probe
grep 'WARP DOWN' output.log    # Container-side tunnel drops
grep 'WARP SHUTDOWN' output.log # Auto-recovery triggered
grep 'WARP RECOVERED' output.log # Self-healed before threshold
```

6 toggle.sh

```bash
#!/bin/bash
# toggle.sh  nullexit gateway toggle script (macOS + Linux)
# Automatically handles Docker containers, Colima, DNS hijacking, and Tailscale exit node routing.
#
# One file, both OSes: the dispatch below runs the self-contained Linux
# implementation (shuf / ss / nmcli / ip / systemd) and exits; on macOS it
# falls through to the macOS implementation (jot / netstat / networksetup,
# launchd / pf).

set -e

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"

#
# LINUX TOGGLE  self-contained; runs the full toggle then exits before the
# macOS body below. Native Linux tooling (shuf / ss / nmcli / ip / systemd).
# (Body inlined verbatim at column 0  it contains multiline `bash -c ""`
# strings that must NOT be re-indented.)
#
if [[ "$OSTYPE" != "darwin"* ]]; then
# Define log file (used by the lifecycle breadcrumb + run_logged helper)
LOG_FILE="$PWD/output.log"

# --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
# TOGGLE_PHASE is updated as the script moves through STOP/START so that if the
# process is killed (terminal closed, pkill, OOM) or exits on error, the EXIT
# trap records the final exit code AND the phase it died in  making an
# interrupted START traceable in output.log instead of ending abruptly.
TOGGLE_PHASE="init"
_log_exit_breadcrumb() {
  local rc=$?
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
    2>/dev/null || true
}
trap '_log_exit_breadcrumb' EXIT

# --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
```

```
# When enabled, mirror a full command-by-command xtrace to output.log. Off by
# default. Read directly from .env because common.sh (read_env_var) isn't
# sourced yet. Branch on bash version: BASH_XTRACEFD needs >= 4.1 (Linux
# usually has bash 5); fall back to teeing stderr on older bash. We never use
# the `exec {var}>>` dynamic-fd form (4.1+ only).
DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'" | tr '[:upper:]' '[:lower:]' | tr -d '[:space:]')
if [ "$DEBUG_TRACE" = "true" ]; then
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled  xtrace mirrored to output.log (bash ${BASH_VERSINFO[0]}.${BASH_VERSINFO[1]})" >> "$LOG_FILE"
  export PS4='+ [${SECONDS}s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: '  # subshell-free clock: a $() in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
  if [ "${BASH_VERSINFO[0]}" -gt 4 ] || { [ "${BASH_VERSINFO[0]}" -eq 4 ] && [ "${BASH_VERSINFO[1]}" -ge 1 ]; }; then
    exec 9>>"$LOG_FILE"
    export BASH_XTRACEFD=9
    set -x
  else
    # bash < 4.1: plain stderrfile redirect. NOT a process substitution  under
    # set -e, xtrace flowing through a backgrounded `tee` can abort the script
    # (failed write/SIGPIPE trips set -e). See devref 15.11.8.
    exec 2>>"$LOG_FILE"
    set -x
  fi
fi

# --- PROCEDURAL PORT GENERATION ---
# To prevent static port fingerprinting by malware, we randomize exposed ports
# on every boot and persist them to a hidden environment file.
PORTS_FILE="/tmp/nullexit-ports.env"
if [[ "$1" == "" || "$1" == "--restart" || "$1" == "restart" ]]; then

  # Collision-avoidance function
  GENERATED_PORTS=""
  get_free_port() {
    local port
    while true; do
      # Linux uses shuf or $RANDOM instead of jot
      port=$(shuf -i 10000-65000 -n 1 2>/dev/null || echo $((10000 + RANDOM % 55000)))
      # 1. Prevent self-collision within this loop
      if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
        continue
      fi
      # 2. Prevent collision with ANY process on the machine (root daemons, terminal apps, etc)
      # We use ss to read the kernel socket table directly.
      if ! ss -lnt | awk '{print $4}' | grep -q ":$port$"; then
        echo "$port"
        return
      fi
    done
  }

  export TOR_SOCKS_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_SOCKS_PORT"

  export TOR_CONTROL_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"

  export SOCKS_PROXY_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"

  export DNS_PROXY_PORT=$(get_free_port)

  echo "export TOR_SOCKS_PORT=$TOR_SOCKS_PORT" > "$PORTS_FILE"
  echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
  echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
  echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
elif [ -f "$PORTS_FILE" ]; then
  # On STOP, source the existing ports so docker-compose down matches
  source "$PORTS_FILE"
else
  # Fallbacks if file is lost
  export TOR_SOCKS_PORT=9050
  export TOR_CONTROL_PORT=9051
  export SOCKS_PROXY_PORT=1080
  export DNS_PROXY_PORT=5354
fi

# Start execution timer
TOGGLE_START_TIME=$SECONDS

```

```
if [[ "$1" == "restart" ]] || [[ "$1" == "--restart" ]]; then
  echo "Executing Gateway Restart Sequence..."
  # We must source common.sh here temporarily to check gateway state before we proceed
  source "$SCRIPT_DIR/scripts/common.sh"
  # Use persisted intent via gateway_state (echoes a string, never a return code;
  # set -e-safe under set -x  see devref 15.11.8). On Linux the marker is not
  # maintained, so this degrades to the is_gateway_active fallback (fast  native Docker).
  if [ "$(gateway_state)" = "running" ]; then
    echo "Gateway is currently running. Stopping it first..."
    bash "$0"
    echo "Gateway stopped. Now starting it..."
  else
    echo "Gateway is already stopped. Starting it..."
  fi
  bash "$0"
  exit 0
fi

# Global flag to track if the script completed successfully
SUCCESS_RUN=false

# Global variable to track the currently running background command PID
CURRENT_BG_PID=""

# Source common bash functions
source "$SCRIPT_DIR/scripts/common.sh"


PID_FILE="/tmp/nullexit-caffeinate.pid"

# Start macOS caffeinate to prevent sleep while gateway is running
start_sleep_prevention() {
  if ! command -v systemd-inhibit >/dev/null 2>&1; then
    echo "  [Warning] systemd-inhibit tool not found. Sleep prevention unavailable."
    return 1
  fi

  # Stop any existing sleep prevention process first to avoid duplicates
  stop_sleep_prevention

  echo "  Enabling system sleep prevention (systemd-inhibit)..."
  # Run systemd-inhibit wrapped in a bash trap to prevent idle system sleep.
  # Critically, this traps shutdown signals (SIGTERM) and automatically
  # flushes hijacked DNS back to normal right before the PC powers off.
  # We do not use recover.sh here because it is a heavy recovery script
  # which takes too long and gets terminated before it can reset the DNS. Instead,
  # we surgically and instantly restore normal DNS settings here.
  nohup bash -c "
    trap '
      echo \"[Shutdown Trap] Reverting network settings...\" >> \"$PWD/output.log\" 2>&1
      kill \$! 2>/dev/null || true
      sudo -n resolvectl domain \"$ACTIVE_SERVICE\" \"\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n resolvectl revert \"$ACTIVE_SERVICE\" >> \"$PWD/output.log\" 2>&1 || true
      resolvectl domain \"$ACTIVE_SERVICE\" \"\" >> \"$PWD/output.log\" 2>&1 || true
      resolvectl revert \"$ACTIVE_SERVICE\" >> \"$PWD/output.log\" 2>&1 || true
      if command -v tailscale >/dev/null 2>&1; then
        tailscale up --accept-dns=false --exit-node= >> \"$PWD/output.log\" 2>&1 &
        TS_DOWN_ARGS=\"\"
        if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
        tailscale down \$TS_DOWN_ARGS >> \"$PWD/output.log\" 2>&1 &
      fi
      sleep 1
      echo \"[Shutdown Trap] Cleanup complete.\" >> \"$PWD/output.log\" 2>&1
      exit 0
    ' SIGTERM SIGINT SIGHUP
    systemd-inhibit --what=sleep --who=nullexit --why=\"gateway running\" --mode=block sleep infinity &
    wait \$!
  " >> output.log 2>&1 &
  local caffe_pid=$!
  echo "$caffe_pid" > "$PID_FILE"
  echo "  Sleep prevention active (PID $caffe_pid). Your PC won't sleep while the gateway is running."
}

# Stop sleep prevention when gateway is stopped
stop_sleep_prevention() {
  if [ -f "$PID_FILE" ]; then
    local caffe_pid
    caffe_pid=$(cat "$PID_FILE")
    if [ -n "$caffe_pid" ] && kill -0 "$caffe_pid" 2>/dev/null && ps -p "$caffe_pid" -o command= 2>/dev/null |
     grep -q -E "systemd-inhibit|recover.sh"; then
      echo "  Stopping system sleep prevention (systemd-inhibit wrapper PID $caffe_pid)..."
```

```
      kill "$caffe_pid" 2>/dev/null || true
    fi
    rm -f "$PID_FILE"
  fi
}

DNS_WATCHER_PID_FILE="/tmp/nullexit-dns-watcher.pid"

stop_dns_watcher() {
  if [ -f "$DNS_WATCHER_PID_FILE" ]; then
    local watcher_pid
    watcher_pid=$(cat "$DNS_WATCHER_PID_FILE")
    if [ -n "$watcher_pid" ] && kill -0 "$watcher_pid" 2>/dev/null; then
      echo "  Stopping background DNS Watcher (PID $watcher_pid)..."
      kill -9 "$watcher_pid" 2>/dev/null || true
    fi
    rm -f "$DNS_WATCHER_PID_FILE"
  fi
}

# Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
cleanup_handler() {
  local exit_code=$?
  local trigger_type="$1"
  local line_no="$2"

  if [ "$SUCCESS_RUN" = "false" ]; then
    echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1...
    "
    if [ -n "$ACTIVE_SERVICE" ]; then
      reset_dns
    fi

    if [ -n "$CURRENT_BG_PID" ]; then
      echo "Terminating active command (PID $CURRENT_BG_PID)..."
      kill -15 "$CURRENT_BG_PID" 2>> output.log || true
    fi

    # Disconnect host Tailscale and close the GUI application on failure
    disconnect_tailscale_host

    # Stop local DNS proxy if running
    stop_local_dns_proxy

    # Stop sleep prevention
    stop_sleep_prevention
    stop_dns_watcher

    # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
    if [ -n "$ACTIVE_SERVICE" ]; then
      cleanup_network_state
    fi


    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n=============================================="
      echo "ERROR: Script failed at line $line_no with exit code $exit_code."
      echo "=============================================="
      echo "Troubleshooting steps:"
      echo "1. Run 'colima status' to verify the VM is active."
      echo "2. Run 'docker compose ps' to check container health."
      echo "3. Run 'docker compose logs' to view application logs."
      echo "4. Run 'tailscale status' to check host Tailscale status."
      echo "=============================================="
    fi
  fi
}

trap 'cleanup_handler ERR $LINENO' ERR
trap 'cleanup_handler INT' INT
trap 'cleanup_handler TERM' TERM
trap 'cleanup_handler HUP' HUP

#  NOPASSWD sudo
# The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
# needed (resolvectl, ip, systemctl, systemd-inhibit, pkill, and
# python3 scripts/dns-proxy.py for the fallback DNS proxy).
# All privileged calls below use `sudo -n` which will run without prompting.
# No credential caching loop is needed.

# Add Homebrew and standard paths
```

```
export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"

# Check for local DNS proxy tools (used when tailnet data plane is unavailable)
# Uses a Python DNS proxy (handles TCP wire format correctly).
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_BIN=""
if command -v python3 >/dev/null 2>&1; then
  DNS_PROXY_BIN="python3 -I"
elif command -v python >/dev/null 2>&1; then
  DNS_PROXY_BIN="python -I"
fi

# Path to the DNS proxy script (sibling of this script)
DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/dns-proxy.py"

# Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
# (`brew install tailscale` + `systemctl start tailscaled`)  NOT the
# Tailscale.app GUI  so there is no .app bundle or Network Extension sandbox
# to launch, no menu-bar click to perform, and no scutil Network Service to
# start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
# if you ran the install with sudo) and the control socket is always available.
TS_BIN=""
if command -v tailscale >> output.log 2>&1; then
  TS_BIN="tailscale"
fi

# Baseline: disable Tailscale DNS management at the daemon level.
# `tailscale set` writes a persistent preference. However, `tailscale up --reset`
# (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
# re-enable `--accept-dns=true` and nullify this `set`.
# Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
# `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
#
# This initial `set` is still useful as a baseline for non-reset operations
# (e.g. if the user runs `tailscale up` manually without --reset) and covers
# the brief window between this line and the first `--reset` call.
#
# Runs once at script start while tailscaled is still connected (if it was
# left running from a previous toggle), so the pref takes effect before any
# disconnection or reconnection logic.
if [ -n "$TS_BIN" ]; then
  $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
fi

# Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
# *entire* teardown  no need to mess with system network managers.
# tailscaled keeps running (as a systemd service), which is intentional:
# the next `tailscale up` then reconnects in seconds rather than waiting for a
# slow daemon launch. The `tailscale down` call also retains the user's
# auth state, so we don't force a browser re-login every cycle.
#
# Defined early (before the if/else branches and referenced by cleanup_handler's
# ERR trap) so that any failure before the main logic can still tear down Tailscale.
restart_tailscaled_daemon() {
  echo "  [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
  echo "  [Tailscale Recovery] Attempting to restart tailscaled service via systemctl..."
  if run_with_timeout 15 sudo -n systemctl restart tailscaled >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled service."
  else
    echo "  [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
  fi
  sleep 3
}
disconnect_tailscale_host() {
  if [ -n "$TS_BIN" ]; then
    echo "Disconnecting host Tailscale from mesh (tailscaled stays running as system service)..."
    local ts_args=""
    if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
    if ! run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1; then
      restart_tailscaled_daemon
      run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1 || true
    fi
  fi
}

ACTIVE_SERVICE=$(get_active_service)

# Helper to nuke leftover network state after teardown (proxy settings, routing
# table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
# problem where stale state kills internet after the gateway stops.
cleanup_network_state() {
```

```
  echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."

  # 1. Disable any leftover proxy settings (needs root on many macOS versions)
  # (Skipped for Linux since we don't set global SOCKS proxy)
  echo "  Proxies disabled."

  # 2. Flush DNS cache
  if command -v resolvectl >> output.log 2>&1; then
    sudo -n resolvectl flush-caches 2>> output.log || true
  elif command -v systemd-resolve >> output.log 2>&1; then
    sudo -n systemd-resolve --flush-caches 2>> output.log || true
  fi
  echo "  DNS cache flushed."

  # 3. Flush stale routing table entries
  if command -v ip >> output.log 2>&1; then
    sudo -n ip route flush cache >> output.log 2>&1 || true
  fi
  echo "  Routing table flushed."

  # 4. Power-cycle Wi-Fi to clear any lingering interface state
  echo "  Restarting network interface..."
  sudo -n ip link set dev "$ACTIVE_SERVICE" down 2>> output.log || true
  sleep 1
  sudo -n ip link set dev "$ACTIVE_SERVICE" up 2>> output.log || true
  echo "  Interface power-cycled ($ACTIVE_SERVICE)."
  sleep 3

  echo "Network state cleanup complete."
}

# Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
# during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
# `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
# would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
# Capture current DNS before resetting so we can check if it was hijacked
INITIAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head
    -1 || true)

echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
reset_dns


  #  Local Network Surveillance Check
  echo -e "\n"
  echo "Local Network Surveillance Check"
  echo ""
  # Count populated ARP cache entries to detect network scanning or high congestion
  if command -v ip >/dev/null 2>&1; then
    ARP_COUNT=$(ip neigh | grep -iv 'incomplete' | wc -l | tr -d ' ')
  else
    # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
    ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
  fi
  if [ "$ARP_COUNT" -gt 15 ]; then
    echo -e "  \033[1;33m[!] WARNING: High ARP activity detected ($ARP_COUNT devices in cache).\033[0m"
    echo -e "  \033[1;33m[!] This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)
     .\033[0m"
    echo -e "  [] Your traffic remains fully encrypted and invisible.\n"
  else
    echo -e "  [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
  fi

echo "Checking Gateway Status..."
# Decide STOP vs START from persisted intent via gateway_state, which echoes a
# string (never a return code) to stay set -e-safe under set -x. See devref 15.11.8.
# On Linux this falls back to is_gateway_active (fast  native Docker).
if [ "$(gateway_state)" = "running" ]; then
  TOGGLE_PHASE="stop"
  echo -e "\n============================================="
  echo "Gateway is RUNNING. Stopping it now..."
  echo -e "=============================================\n"

  # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
  disconnect_tailscale_host
  echo ""

  echo "Stopping Docker containers..."
  # `|| true`: a disable must always complete even if the Docker daemon is down.
  run_logged docker compose down -t 5 || true

```

```
  echo -e "\nDocker daemon handles container teardown automatically."

  # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
  # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
  # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
  echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
  reset_dns

  # 3. Stop local DNS proxy if running
  stop_local_dns_proxy

  # Stop sleep prevention
  stop_sleep_prevention
  stop_dns_watcher

  # 4. Nuke leftover network state so internet actually works after teardown
  cleanup_network_state

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
else
  TOGGLE_PHASE="start"
  echo -e "\n============================================="
  echo "Gateway is STOPPED. Starting it now..."
  echo -e "=============================================\n"

  # 1. Prevent host exit-node deadlock during VM / Container startup
  disconnect_tailscale_host

  echo -e "\nUsing native Linux Docker..."

  # 4. Clean up corrupted AdGuardHome configurations
  if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
    rm -f "adguard/conf/AdGuardHome.yaml"
    echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
  fi



  # 5. Start compose services
  write_host_ips
  # Procedurally generated ports have already been exported at the top of the script.
  echo "  [Security] Procedurally randomized proxy ports generated to evade fingerprinting."

  # --- HONEY-PORT TRIPWIRE ---
  # Generate a random trap port. If any malware does a sequential port scan on localhost,
  # it will inevitably hit this port. When it does, we instantly fire a desktop alert.
  HONEY_PORT=$(get_free_port)
  (
    if nc -l -p "$HONEY_PORT" 127.0.0.1 >/dev/null 2>&1 || nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$PWD/output.log"
      if command -v notify-send >/dev/null 2>&1; then
        notify-send -u critical "nullexit OPSEC Alert" "An unknown application is aggressively scanning your local ports. Malware port-scan suspected."
      fi
    fi
  ) &
  disown %1 2>/dev/null || true
  echo "  [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."

  echo -e "\nStarting Docker containers..."
  run_logged docker compose up -d

  # Clean up the one-off rule compiler container so it doesn't clutter the Docker UI
  docker compose rm -s -f rule-compiler >> output.log 2>&1

  # 6. Wait for the gateway container's Tailscale connection to be ready
  BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
  USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')

  echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
  echo "  (This can take 30-60s  Tailscale goes through: NoState  Starting  Running  Online)"
  connected=false
  consecutive_failures=0
  for i in {1..60}; do
    # Run tailscale status and capture its output so the user can see what's happening.
    # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
    ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)

    if echo "$ts_output" | grep -q "offers exit node"; then
```

```
      connected=true
      break
    fi

    # Track consecutive failures to detect a stuck state.
    # If we see 40+ consecutive 'NoState' responses, give up early
    # the daemon is alive but can't connect to the coordination server.
    if echo "$ts_output" | grep -q "NoState"; then
      consecutive_failures=$((consecutive_failures + 1))
      if [ "$consecutive_failures" -ge 40 ]; then
        echo ""
        echo "  [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
        echo "    Tailscale daemon is running but can't reach the coordination server."
        break
      fi
    else
      consecutive_failures=0
    fi

    # Print a condensed status every 10 seconds so the user can see progress
    if [ $((i % 10)) -eq 0 ]; then
      ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
      echo "  [attempt $i/60] $ts_short"
    elif [ $((i % 5)) -eq 0 ]; then
      echo -n "[$i]"
    else
      echo -n "."
    fi
    sleep 1
  done
  echo ""

  if [ "$connected" = "true" ]; then
    echo "Gateway container is online on the Tailscale mesh."
  elif [ "$BYPASS_PING" = "true" ]; then
    echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
  else
    echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
    echo "  The container was seen in the following states during the wait:" >&2
    echo "$ts_output" | sed 's/^/    /' >&2
    echo "" >&2
    echo "  This could mean:" >&2
    echo "    1. Tailscale auth key has expired  check TS_AUTHKEY in .env" >&2
    echo "    2. Network issue inside the container  check: docker compose logs tailscale" >&2
    echo "    3. gluetun VPN tunnel not connecting  check: docker compose logs warp | tail -20" >&2
    echo "" >&2
    echo "  Diagnose manually:" >&2
    echo "    docker compose exec -T tailscale tailscale status" >&2
    echo "    docker compose exec -T tailscale tailscale netcheck" >&2
    cleanup_handler ERR $LINENO
    exit 1
  fi

  # 6b. Resolve the gateway's Tailscale IP.
  # Dynamically query the container for the IP (handles IP changes after re-auth)
  # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
  #
  # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
  # returns (\r) into captured output. If $TS_IP ends with \r, then
  # `networksetup -setdnsservers` silently rejects the malformed IP and
  # DNS stays at 1.1.1.1  the primary bug this project was hitting.
  # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
  # the first field of the first line.
  TS_IP=""
  if [ "$connected" = "true" ]; then
    TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
     awk 'NR==1{print $1; exit}' || true)
    if [ -n "$TS_IP" ]; then
      echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
      echo "$TS_IP" > .gateway_ip
    else
      echo "    [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
      TS_IP=$(read_adguard_ip || true)
    fi
  else
    TS_IP=$(read_adguard_ip || true)
  fi

  # 7. Connect host Mac to Tailscale mesh.
  # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
  # system LaunchDaemon via 'systemctl start tailscaled'), the previous five-
```

```
# phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
# no menu bar item to click, and no Accessibility permission required.
#
# CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
# We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
# if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
# are impossible without the host on the mesh).
#
# We connect in two phases:
#   Phase A  Join mesh WITHOUT exit node (so pre-flight checks can run)
#   Phase B  Set exit node only after pre-flight checks confirm the gateway works
HOST_ON_MESH=false
SKIP_EXIT_NODE=false
if [ -n "$TS_BIN" ]; then
  echo -n "Verifying tailscaled is reachable"
  daemon_ready=false
  for i in {1..10}; do
    # `tailscale status` returns exit code 1 when the daemon is alive but
    # disconnected ("Tailscale is stopped."). We check the daemon socket
    # directly as a fallback  if the socket exists, tailscaled is running
    # even if the host isn't connected to the mesh yet.
    if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
      daemon_ready=true
      break
    fi
    echo -n "."
    sleep 1
  done
  echo ""

  if [ "$daemon_ready" != "true" ]; then
    echo -e "\033[0;33m[!] tailscaled daemon is not running. Attempting to start it...\033[0m"

    # Try system-wide daemon (with timeout  sudo may prompt for password)
    if [ "$daemon_ready" != "true" ]; then
      if run_with_timeout 10 sudo systemctl start tailscaled 2>> output.log; then
        echo "  Started tailscaled as system LaunchDaemon."
        for i in {1..15}; do
          if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
            daemon_ready=true
            break
          fi
          sleep 1
        done
      fi
    fi
  fi

  if [ "$daemon_ready" = "true" ]; then
    echo "tailscaled is responsive (running as a system service)."

    #  Phase A: Join mesh without exit node
    echo "Connecting host to Tailscale mesh (no exit node yet)..."
    if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node=; then
      HOST_ON_MESH=true
      echo "Host is on Tailscale mesh."
    else
      echo -e "\033[0;33m[!] tailscale up failed even though the daemon is running.\033[0m"
      echo "  This usually means the host hasn't authenticated with your tailnet yet."
      echo "  Run this command in a terminal to authenticate:"
      echo ""
      echo "      sudo tailscale up"
      echo ""
      echo "  A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
    fi
  else
    echo -e "\033[0;31m[!] tailscaled could NOT be started automatically.\033[0m"
    echo "  Run this in a terminal to start and authenticate Tailscale:"
    echo ""
    echo "      sudo systemctl start tailscaled"
    echo "      sudo tailscale up"
    echo ""
    echo "  Then re-run toggle.sh."
  fi
else
  echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
fi

#  Phase B: Pre-flight checks + enable exit node only if safe
if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
  echo -e "\n"
```

```
    echo "Pre-flight connectivity check"
    echo ""

    # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
    # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
    # ("direct connection not established") even though the pong was received.
    # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
    echo -n "  [1/3] Gateway reachable via Tailscale... "
    ping_ok=false
    # The host tailscaled needs a few seconds to propagate the new network map and
    # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
    for attempt in {1..45}; do
      if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
        ping_ok=true
        break
      fi
      sleep 1
    done
    if [ "$ping_ok" = "true" ]; then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
    echo -n "  [2/3] AdGuard DNS via localhost... "
    if command -v dig >/dev/null 2>&1; then
      if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 >/dev/null 2>&1; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    elif command -v nslookup >/dev/null 2>&1; then
      if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 >/dev/null 2>&1; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    else
      echo "SKIP (dig/nslookup not available)"
    fi

    # Check 3: Does the WARP container have external internet?
    echo -n "  [3/3] WARP container internet... "
    if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace &>/dev/null;
     then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Act based on check results
    if [ "$SKIP_EXIT_NODE" != "true" ]; then
      echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
      if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node="$TS_IP" --exit-node-allow-lan-access=
     true; then
        echo "Exit node enabled."
      else
        echo "[Warning] Failed to set exit node."
        SKIP_EXIT_NODE=true
      fi
    fi

    if [ "$SKIP_EXIT_NODE" = "true" ]; then
      echo -e "\n\033[0;33m[!] Pre-flight checks failed  EXIT NODE + DNS HIJACK SKIPPED.\033[0m"
      echo "  Host stays on Tailscale mesh but your internet routes through your normal connection."
      echo "  Troubleshoot with:"
      echo "    docker compose logs warp | tail -20"
      echo "    docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"
      echo ""
      echo "  Falling back to SOCKS5 proxy for traffic routing..."
    fi
  elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
    echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
    SKIP_EXIT_NODE=true
  fi

```

```
  # 8. Apply DNS Hijacking.
  # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
  # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
  # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
  if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
    echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
    echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
    if force_dns_to_gateway "$TS_IP"; then
      echo "DNS hijack verified: single server = $TS_IP."
      start_dns_watcher "$TS_IP"
    else
      echo "[Warning] DNS hijack could not be verified."
      echo "  If DNS is broken, run manually:"
      echo "    networksetup -setdnsservers \"$ACTIVE_SERVICE\" $TS_IP"
      echo "    networksetup -setsearchdomains \"$ACTIVE_SERVICE\" ts.net"
    fi
  elif [ -n "$TS_IP" ]; then
    echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
    echo "  Trying local DNS proxy for ad-blocking..."
    if start_local_dns_proxy; then
      echo "  Ad-blocking active via local DNS proxy through AdGuard."
    else
      echo "  Local DNS proxy unavailable. DNS stays at 1.1.1.1."
      echo "  AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
    fi

    # Enable SOCKS5 proxy for traffic routing through WARP
    # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
    echo -e "\n  Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
    echo "  (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
    enable_socks_proxy "$ACTIVE_SERVICE"

    # Verify the SOCKS5 proxy works through WARP
    echo -n "  Verifying SOCKS5 proxy routes through WARP... "
    if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn-cgi/
     trace 2>/dev/null | grep -q "warp=on"; then
      echo "ok (warp=on)."
      echo "  All TCP traffic now encrypted through Cloudflare WARP tunnel."
    else
      echo "FAILED (proxy not routing through WARP)."
      echo "  Disabling SOCKS5 proxy  internet stays on direct connection."
      disable_socks_proxy
      echo "  SOCKS proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
    fi
  else
    echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
  fi

  # 9. Verify host connectivity
  if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
    if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
      echo "Host is online on the Tailscale mesh."
    else
      echo "[Warning] Host is not responding on Tailscale."
    fi
  fi

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."

  # Prevent system from going to sleep while gateway is running
  start_sleep_prevention
fi

SUCCESS_RUN=true

# Final DNS state summary
if [ -n "$ACTIVE_SERVICE" ]; then
  FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr '
' ' ' || true)
  echo -e "\n"
  echo -e "DNS STATE: $FINAL_DNS"
  echo -e ""
fi

if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
  if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
    echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
    echo "  The gateway gracefully fell back to yesterday's cached blocklist."
    echo "  (Run 'docker logs rule-compiler' later to investigate which link died.)"
  fi
```

```
fi

echo -e "\n"
read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
  echo "Reversing gateway state..."
  exec bash "$0"
fi

echo -e "\nYou can close this terminal window now."

exit 0
fi

#
# macOS TOGGLE  jot / netstat / networksetup, launchd / pf.
#

# Enforce Cryptographic Script Integrity
if [ -f "scripts/crypto.sh" ]; then
  if ! bash scripts/crypto.sh --verify; then
    exit 1
  fi
fi

# Define log file
LOG_FILE="$PWD/output.log"

# Rotate log file if it exceeds 1GB (1073741824 bytes).
# Sized large on purpose: with DEBUG_TRACE the log holds a full command-by-command
# trace  effectively the entire compilation/warning history of the framework
# which is valuable to keep for post-mortems. Raised from 50MB so always-on
# tracing doesn't churn that history away prematurely.
if [ -f "$LOG_FILE" ]; then
  log_size=$(stat -f%z "$LOG_FILE" 2>/dev/null || stat -c%s "$LOG_FILE" 2>/dev/null || echo 0)
  if [ "$log_size" -gt 1073741824 ]; then
    mv "$LOG_FILE" "${LOG_FILE}.old" 2>/dev/null || rm -f "$LOG_FILE"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] Log file exceeded 1GB; rotated to output.log.old" > "$LOG_FILE"
  fi
fi

# --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
# TOGGLE_PHASE is updated as the script moves through STARTUP/STOP/START so that
# if the process is killed (window closed, pkill, OOM) or exits with an error,
# the EXIT trap records the final exit code AND the phase it died in. Previously
# an interrupted START left output.log ending abruptly after the teardown with no
# indication of what failed  this makes every exit traceable.
TOGGLE_PHASE="init"
_log_exit_breadcrumb() {
  local rc=$?
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
    2>/dev/null || true
}
trap '_log_exit_breadcrumb' EXIT

# --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
# When enabled, mirror a full command-by-command xtrace to output.log for a
# complete post-mortem trail (e.g. a race during --restart while Wi-Fi is
# re-associating). Off by default. Read directly from .env here because
# common.sh (read_env_var) isn't sourced yet.
#
# BASH_XTRACEFD (send xtrace to its own fd, keeping the terminal clean) needs
# bash >= 4.1. macOS ships /bin/bash 3.2, so we branch:
#   - bash >= 4.1 (e.g. Linux): dedicate fd 9 to the log and point xtrace there;
#     the terminal stays clean and normal stderr is untouched.
#   - bash 3.2 (stock macOS): BASH_XTRACEFD is unsupported, so xtrace always goes
#     to stderr. We redirect stderr *straight to the log file* with a plain
#     `exec 2>>"$LOG_FILE"`. This deliberately does NOT use a process
#     substitution (`2> >(tee )`): under `set -e`, stderr flowing through the
#     backgrounded `tee` proved to abort the script mid-START on 3.2 (a failed
#     write / SIGPIPE in a pipeline returned non-zero and tripped `set -e`, with
#     $LINENO mis-blaming the enclosing `fi`  see devref 15.11.8-A/E). The
#     tradeoff: on 3.2 the terminal no longer shows live progress while tracing
#     (everything is in output.log instead). Acceptable for an opt-in debug mode.
# We NEVER use the `exec {var}>>` dynamic-fd form (4.1+ only), which hard-fails on 3.2.
DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'" | tr '[:upper:]' '[:lower
    :]' | tr -d '[:space:]')
if [ "$DEBUG_TRACE" = "true" ]; then
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled  xtrace mirrored to output.log (bash ${BASH_VERSINFO
    [0]}.${BASH_VERSINFO[1]})" >> "$LOG_FILE"
  # Timestamped, location-aware xtrace prompt (works on 3.2 and 4.x).
```

```
  export PS4='+ [${SECONDS}s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: '  # subshell-free clock: a $(...) in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
  if [ "${BASH_VERSINFO[0]}" -gt 4 ] || { [ "${BASH_VERSINFO[0]}" -eq 4 ] && [ "${BASH_VERSINFO[1]}" -ge 1 ]; }; then
    exec 9>>"$LOG_FILE"
    export BASH_XTRACEFD=9
    set -x
  else
    # bash 3.2 (stock macOS): plain stderrfile redirect (no process substitution).
    exec 2>>"$LOG_FILE"
    set -x
  fi
fi

# --- MAIN EXECUTION LOGGING ---
echo -e "\n=======================================================" >> "$LOG_FILE"
echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh invoked" >> "$LOG_FILE"
echo "=======================================================" >> "$LOG_FILE"

# (Previously this script chmod 600'd .env at start to "unlock" it for docker compose,
# and chmod 000'd it on exit. That pattern broke docker compose on macOS/Colima when
# the umask produced 0600  the script removed its own ability to read .env before it could even proceed.)

# --- PROCEDURAL PORT GENERATION ---
# To prevent static port fingerprinting by malware, we randomize exposed ports
# on every boot and persist them to a hidden environment file.
PORTS_FILE="/tmp/nullexit-ports.env"
if [[ "$1" == "" || "$1" == "--restart" ]]; then

  # Collision-avoidance function
  GENERATED_PORTS=""
  get_free_port() {
    local port
    while true; do
      port=$(jot -r 1 10000 65000)
      # 1. Prevent self-collision within this loop
      if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
        continue
      fi
      # 2. Prevent collision with ANY process on the Mac (root daemons, terminal apps, etc)
      # We use netstat instead of lsof because netstat reads the kernel socket table
      # directly and doesn't require sudo to see ports bound by other users/root.
      if ! netstat -an -p tcp | awk '{print $4}' | grep -q "\.$port$"; then
        echo "$port"
        return
      fi
    done
  }

  export TOR_SOCKS_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_SOCKS_PORT"

  export TOR_CONTROL_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"

  export SOCKS_PROXY_PORT=$(get_free_port)
  GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"

  export DNS_PROXY_PORT=$(get_free_port)

  echo "export TOR_SOCKS_PORT=$TOR_SOCKS_PORT" > "$PORTS_FILE"
  echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
  echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
  echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
  ADGUARD_PASS=$(grep -E "^ADGUARD_PASSWORD=" .env 2>/dev/null | cut -d'=' -f2- | tr -d "\"'")
  echo "export TOR_PASSWORD=${ADGUARD_PASS:-nullexit}" >> "$PORTS_FILE"
elif [ -f "$PORTS_FILE" ]; then
  # On STOP, source the existing ports so docker-compose down matches
  source "$PORTS_FILE"
else
  # Fallbacks if file is lost
  export TOR_SOCKS_PORT=9050
  export TOR_CONTROL_PORT=9051
  export SOCKS_PROXY_PORT=1080
  export DNS_PROXY_PORT=5354
fi

# Start execution timer
TOGGLE_START_TIME=$SECONDS

if [[ "$1" == "--restart" ]]; then
```

```
  echo "Executing Gateway Restart Sequence..."
  # We must source common.sh here temporarily to check gateway state before we proceed
  source "$SCRIPT_DIR/scripts/common.sh"
  # Use persisted intent (marker) via gateway_state, which echoes a string (never
  # a return code) to stay set -e-safe under set -x. See devref 15.11.8.
  if [ "$(gateway_state)" = "running" ]; then
    echo "Gateway is currently running. Stopping it first..."
    bash "$0"
    echo "Gateway stopped. Now starting it..."
  else
    echo "Gateway is already stopped. Starting it..."
  fi
  bash "$0"
  exit 0
fi

# Global flag to track if the script completed successfully
SUCCESS_RUN=false

# Global variable to track the currently running background command PID
CURRENT_BG_PID=""

# Source common bash functions
source "$SCRIPT_DIR/scripts/common.sh"

# Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
LOCK_FILE="/tmp/nullexit-toggle.lock"
if [ -f "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
      die "Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running."
    fi
  fi
fi
echo "$$" > "$LOCK_FILE"

# Helper function to run a GUI command as the logged-in console user
# (Prevents permission failures if the script is run with sudo)
run_gui_cmd() {
  local console_user
  console_user=$(stat -f '%Su' /dev/console 2>> output.log || echo "$SUDO_USER")
  if [ -z "$console_user" ] || [ "$console_user" = "root" ]; then
    console_user=$(logname 2>> output.log || echo "$USER")
  fi

  if [ -n "$console_user" ] && [ "$console_user" != "root" ] && [ "$EUID" -eq 0 ]; then
    sudo -u "$console_user" "$@"
  else
    "$@"
  fi
}

# Active-state marker: written at end of START path so external watchers
# (see launchd/com.nullexit.wake-recovery.plist + scripts/watcher.sh +
# devref 10.29) know the gateway is currently up. They only fire
# recover.sh --post-wake when this file is present. Cleared at the top
# of the STOP path and inside cleanup_handler on any error/signal path
# so a half-broken gateway never re-triggers post-wake refreshes.
write_gateway_active_marker() {
  date -u +%FT%TZ > /tmp/nullexit-gateway-active.marker
  # Set the watcher debounce timestamp to now to prevent a redundant post-startup recovery run.
  date +%s > /tmp/nullexit-watcher.last-recovery 2>/dev/null || true
}

clear_gateway_active_marker() {
  rm -f /tmp/nullexit-gateway-active.marker
  rm -f "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/TUNNEL_FAILED_CLOSED.marker"
}

# Capture whether the gateway-active marker exists BEFORE the defensive clear
# below wipes it. The STOP-vs-START decision (gateway_should_be_running) reads
# this captured value  persisted *intent*  instead of a live Docker probe, so
# a disable is instant and works even when Colima/Docker is down. Must be read
# before clear_gateway_active_marker or it would always look "stopped".
#
# Written as default-assign + `if` WITHOUT an `else`: under `set -e` (with the
# DEBUG_TRACE xtrace redirect active) an `if [ -f  ]; then ; else ; fi` whose
# condition is false was aborting the script here at init on macOS /bin/bash 3.2
#  the false `[ -f ]` status leaked through the compound. An `if` with no `else`
# yields status 0 when the condition is false, so nothing leaks.
```

```
GATEWAY_MARKER_AT_STARTUP="no"
if [ -f "$MARKER_FILE" ]; then
  GATEWAY_MARKER_AT_STARTUP="yes"
fi

# Defensive: clear any stale marker from a prior crashed/aborted run. If a
# previous toggle.sh wrote the marker but was OOM-killed in the middle of the
# START block (before the END-of-START write)  or if the user rebooted mid
# session  this prevents the watcher from firing post-wake against a
# non-running gateway on every wake event. Placed AFTER the function
# definitions so bash resolves the call at parse time.
clear_gateway_active_marker

PID_FILE="$PID_CAFFEINATE"

# Start macOS caffeinate to prevent sleep while gateway is running
start_sleep_prevention() {
  if ! command -v caffeinate >/dev/null 2>&1; then
    echo "  [Warning] caffeinate tool not found. Sleep prevention unavailable."
    return 1
  fi

  # Stop any existing sleep prevention process first to avoid duplicates
  stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"

  echo "  Enabling system sleep prevention (caffeinate)..."
  # Run caffeinate wrapped in a bash trap to prevent idle system sleep.
  # Critically, this traps macOS shutdown signals (SIGTERM) and automatically
  # flushes hijacked DNS back to normal right before the Mac powers off.
  # We do not use recover.sh here because it is a heavy recovery script (managing
  # containers, restarting tailscaled, etc.) which takes too long and gets
  # terminated by macOS before it can reset the DNS. Instead, we surgically and
  # instantly restore normal DNS and proxy settings here.
  nohup bash -c "
    trap '
      echo \"[Shutdown Trap] Reverting network settings...\" >> \"$PWD/output.log\" 2>&1
      kill \$! 2>/dev/null || true
      sudo -n networksetup -setdnsservers \"$ACTIVE_SERVICE\" 1.1.1.1 >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setdnsservers \"$ENO_SERVICE\" 1.1.1.1 >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsearchdomains \"$ACTIVE_SERVICE\" \"Empty\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsearchdomains \"$ENO_SERVICE\" \"Empty\" >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsocksfirewallproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 ||
    true
      sudo -n networksetup -setsocksfirewallproxystate \"$ENO_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setwebproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setwebproxystate \"$ENO_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsecurewebproxystate \"$ACTIVE_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      sudo -n networksetup -setsecurewebproxystate \"$ENO_SERVICE\" off >> \"$PWD/output.log\" 2>&1 || true
      if command -v tailscale >/dev/null 2>&1; then
        tailscale up >> \"$PWD/output.log\" 2>&1 || true
        tailscale set --accept-dns=false --exit-node= >> \"$PWD/output.log\" 2>&1 || true
        TS_DOWN_ARGS=\"\"
        if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
        tailscale down \$TS_DOWN_ARGS >> \"$PWD/output.log\" 2>&1 &
      fi
      sleep 1
      echo \"[Shutdown Trap] Cleanup complete.\" >> \"$PWD/output.log\" 2>&1
      exit 0
    ' SIGTERM SIGINT SIGHUP
    caffeinate -i &
    wait \$!
  " >> output.log 2>&1 &
  local caffe_pid=$!
  echo "$caffe_pid" > "$PID_FILE"
  echo "  Sleep prevention active (PID $caffe_pid). Your Mac won't sleep while the gateway is running."
}



DNS_WATCHER_PID_FILE="$PID_DNS_WATCHER"
WARP_WATCHER_PID_FILE="/tmp/nullexit-warp-watcher.pid"



# Background WARP tunnel liveness monitor. Polls cdn-cgi/trace every 30s
# while healthy, but accelerates to polling every 5s if a failure is detected.
# Always logs state transitions to output.log. When warp=off persists for
# WARP_FAIL_THRESHOLD consecutive polls (default 6 = 30s of downtime), the watcher
# triggers a nuclear recovery: runs recover.sh to tear down the entire
# gateway  disconnecting Tailscale, stopping containers, resetting DNS,
# and power-cycling Wi-Fi. Note: This minimizes exposure window, but
```

```
# only the pf kill switch guarantees absolutely zero leakage on drop.
#
# The threshold (default 6) is statistically chosen: even at an
# unrealistically high 10% false-positive rate per poll, P(6 consecutive
# false positives) = 0.1^6  0.0001%. Real-world false-positive rates
# are far lower, making 30s of continuous downtime overwhelmingly likely
# to be a genuine outage.
#
# Configurable via .env:  WARP_FAIL_THRESHOLD=3  (15s, more aggressive)
start_warp_watcher() {
  stop_warp_watcher
  # Parse threshold from .env (same pattern as GATEWAY_MSS, GATEWAY_BYPASS_PING, etc.)
  local threshold
  threshold=$(read_env_var WARP_FAIL_THRESHOLD)
  if [ -z "$threshold" ]; then
    threshold="${WARP_FAIL_THRESHOLD:-6}"
  fi
  local trigger_file="/tmp/nullexit-warp-shutdown-triggered"
  rm -f "$trigger_file"
  echo "  Starting background WARP Watcher (polling every 30s [accelerating to 5s on failure], shutdown after ${threshold} consecutive failures)..."
  nohup bash -c "
    trap 'exit 0' SIGTERM SIGINT SIGHUP
    last_state='on'
    consec_off=0
    threshold='$threshold'
    trigger_file='$trigger_file'
    out_log='$PWD/output.log'
    recover_bin='$PWD/recover.sh'
    inhibit_file='/tmp/nullexit-warp-inhibit.marker'
    while true; do
      #  Inhibit check
      # recover.sh --post-wake writes this marker while force-recreating the
      # warp container. During that window, the container is intentionally
      # down  don't count it as a failure or we'll fire nuclear recover.sh
      # and kill a gateway that was in the process of healing itself.
      if [ -f \"\$inhibit_file\" ]; then
        consec_off=0
        sleep 5
        continue
      fi

      state='off'
      if docker compose exec -T warp wget -qO- --timeout=5 \
           https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
           | grep -q 'warp=on'; then
        state='on'
      fi

      #  State-transition logging
      if [ \"\$state\" != \"\$last_state\" ]; then
        if [ \"\$state\" = 'off' ]; then
          echo \"[\$(date -u +%FT%TZ)] WARP DOWN  cdn-cgi/trace reports warp=off\" >> \"\$out_log\"
        fi
        last_state=\"\$state\"
      fi

      #  Consecutive-failure tracking + auto-shutdown
      if [ \"\$state\" = 'off' ]; then
        consec_off=\$((consec_off + 1))
        if [ \"\$consec_off\" -ge \"\$threshold\" ] && [ ! -f \"\$trigger_file\" ]; then
          touch \"\$trigger_file\"
          echo \"[\$(date -u +%FT%TZ)] WARP SHUTDOWN  \$consec_off consecutive failures (threshold=\$threshold). Running recover.sh to kill the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.\" >> \"\$out_log\"
          # Nuclear recovery: tear down the entire gateway.
          # recover.sh resets DNS  1.1.1.1, disconnects Tailscale,
          # stops Docker containers, flushes routes, power-cycles Wi-Fi.
          bash \"\$recover_bin\" --auto >> \"\$out_log\" 2>&1 || true
          echo \"[\$(date -u +%FT%TZ)] WARP SHUTDOWN  recover.sh completed, watcher exiting. Your IP is now your real ISP IP.\" >> \"\$out_log\"
          exit 0
        fi
      else
        if [ \"\$consec_off\" -gt 0 ]; then
          echo \"[\$(date -u +%FT%TZ)] WARP RECOVERED  warp=on after \$consec_off consecutive off readings (threshold was \$threshold)\" >> \"\$out_log\"
        fi
        consec_off=0
      fi

```

```
      if [ \"\$state\" = \"off\" ]; then
        sleep 5
      else
        sleep 30
      fi
    done
  " >> output.log 2>&1 &
  echo $! > "$WARP_WATCHER_PID_FILE"
}

stop_warp_watcher() {
  if [ -f "$WARP_WATCHER_PID_FILE" ]; then
    local wp
    wp=$(cat "$WARP_WATCHER_PID_FILE")
    if [ -n "$wp" ] && kill -0 "$wp" 2>/dev/null; then
      echo "  Stopping background WARP Watcher (PID $wp)..."
      kill "$wp" 2>/dev/null || true
    fi
    rm -f "$WARP_WATCHER_PID_FILE"
  fi
}




# Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
cleanup_handler() {
  local exit_code=$?
  local trigger_type="$1"
  local line_no="$2"

  if [ "$SUCCESS_RUN" = "false" ]; then
    echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1...
    "
    if [ -n "$ACTIVE_SERVICE" ]; then
      reset_dns
    fi

    if [ -n "$CURRENT_BG_PID" ]; then
      echo "Terminating active command (PID $CURRENT_BG_PID)..."
      kill -15 "$CURRENT_BG_PID" 2>> output.log || true
    fi

    # Disconnect host Tailscale and close the GUI application on failure
    disconnect_tailscale_host

    # Stop local DNS proxy if running
    stop_local_dns_proxy

    # Stop sleep prevention
    stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
    stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
    stop_warp_watcher
    clear_gateway_active_marker

    # Capture warp logs on failure for debugging before teardown
    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n--- WARP FAILURE LOGS ---" >> output.log
      docker compose logs warp --tail=100 >> output.log 2>&1 || true
      echo "-------------------------" >> output.log
    fi

    # Best-effort container teardown on failure
    docker compose down --remove-orphans -t 30 2>> output.log || true

    # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
    if [ -n "$ACTIVE_SERVICE" ]; then
      cleanup_network_state
    fi


    if [ "$trigger_type" = "ERR" ]; then
      echo -e "\n==============================================="
      echo "ERROR: Script failed at line $line_no with exit code $exit_code."
      echo "==============================================="
      echo "Troubleshooting steps:"
      echo "1. Run 'colima status' to verify the VM is active."
      echo "2. Run 'docker compose ps' to check container health."
      echo "3. Run 'docker compose logs' to view application logs."
      echo "4. Run 'tailscale status' to check host Tailscale status."
```

```
      echo "============================================="
    fi
    rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
  fi
}

trap 'cleanup_handler ERR $LINENO' ERR
trap 'cleanup_handler INT' INT
trap 'cleanup_handler TERM' TERM
trap 'cleanup_handler HUP' HUP

#  NOPASSWD sudo
# The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
# needed (networksetup, pfctl, route, ifconfig, dscacheutil, killall, pkill,
# and python3 -I scripts/dns-proxy.py for the fallback DNS proxy).
# All privileged calls below use `sudo -n` which will run without prompting.
# No credential caching loop is needed.

# Add Homebrew and standard paths
setup_standard_path

# Check for local DNS proxy tools (used when tailnet data plane is unavailable)
# Uses a Python DNS proxy (handles TCP wire format correctly).
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_BIN=""
if command -v python3 >> output.log 2>&1; then
  DNS_PROXY_BIN="python3 -I"
elif command -v python >> output.log 2>&1; then
  DNS_PROXY_BIN="python -I"
fi

# Path to the DNS proxy script (sibling of this script)
DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/dns-proxy.py"

# Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
# (`brew install tailscale` + `brew services start tailscale`)  NOT the
# Tailscale.app GUI  so there is no .app bundle or Network Extension sandbox
# to launch, no menu-bar click to perform, and no scutil Network Service to
# start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
# if you ran the install with sudo) and the control socket is always available.
TS_BIN=""
if command -v tailscale >> output.log 2>&1; then
  TS_BIN="tailscale"
fi

# Baseline: disable Tailscale DNS management at the daemon level.
# `tailscale set` writes a persistent preference. However, `tailscale up --reset`
# (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
# re-enable `--accept-dns=true` and nullify this `set`.
# Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
# `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
#
# This initial `set` is still useful as a baseline for non-reset operations
# (e.g. if the user runs `tailscale up` manually without --reset) and covers
# the brief window between this line and the first `--reset` call.
#
# Runs once at script start while tailscaled is still connected (if it was
# left running from a previous toggle), so the pref takes effect before any
# disconnection or reconnection logic.
if [ -n "$TS_BIN" ]; then
  $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
fi

# Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
# *entire* teardown  no scutil tunnel stop, no Tailscale.app GUI quit, no pkill.
# tailscaled keeps running (as a LaunchAgent / LaunchDaemon), which is intentional:
# the next `tailscale up` then reconnects in seconds rather than waiting for a
# slow .app launch + Network Extension activation. The `tailscale down` call also
# retains the user's auth state, so we don't force a browser re-login every cycle.
#
# Defined early (before the if/else branches and referenced by cleanup_handler's
# ERR trap) so that any failure before the main logic can still tear down Tailscale.


#  Wait for Network Readiness
# On --restart, the gateway tears down while the host network interface may
# still be reconnecting (Wi-Fi, Ethernet, USB tethering, etc.). If we proceed
# too early, get_active_service returns nothing, routing targets the wrong
# interface, and DNS/WARP setup silently fails.
#
# We use `route get default` to detect a valid default gateway rather than
```

```
# checking a specific interface (e.g. en0)  this generalizes across all
# connection types: Wi-Fi, Ethernet, USB-C adapters, iPhone tethering, etc.
# The moment the OS has any routable default gateway, we proceed.
_net_wait_secs=0
_net_timeout=60
echo "Waiting for network connectivity before starting..."
while true; do
  if route get default 2>/dev/null | grep -q 'gateway:'; then
    _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
    echo "  Network ready (default gateway: $_gw)."
    break
  fi
  if [ "$_net_wait_secs" -ge "$_net_timeout" ]; then
    echo "  [Warning] No default gateway after ${_net_timeout}s  proceeding anyway."
    break
  fi
  sleep 2
  _net_wait_secs=$((_net_wait_secs + 2))
done

ACTIVE_SERVICE=$(get_active_service)

# Resolve the service name for en0 (usually "Wi-Fi")  macOS scutil DNS resolver
# is commonly scoped to en0, so per-service DNS changes on other interfaces are
# ignored unless en0's service is also updated.
EN0_SERVICE=$(get_en0_service)

# Helper to add host-side bypass routes for the WARP WireGuard endpoints.
# This prevents an infinite recursive routing loop (tunnel loop) where
# the container's WireGuard packets exit the VM, reach the host, match the
# default route (the exit node), and get routed back into the tunnel.
# We route via the gateway IP (not interface) to ensure packets are routed
# properly even when the host's default route changes to the Tailscale interface.

# Helper to manually override the host's default route to point through the
# Tailscale utun* interface. Standalone tailscaled on macOS (Homebrew version)
# lacks the platform-specific integration to override the default gateway
# automatically when --exit-node is used.
setup_exit_node_routing() {
  local ts_iface
  ts_iface=$(ifconfig | awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':')

  if [ -n "$ts_iface" ]; then
    echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
    local host_mtu
    host_mtu=$(read_env_var HOST_MTU)
    host_mtu=${host_mtu:-1200}

    # Lower the host Tailscale MTU (defaults to 1200) to prevent fragmentation when passing
    # through the container's WARP tunnel (which also has an MTU of 1280).
    sudo -n ifconfig "$ts_iface" mtu "$host_mtu" >> output.log 2>&1 || true

    # DO NOT delete the physical default route! It crashes Colima.
    # Use the /1 trick to mathematically override the default route.
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true

    if netstat -nr | grep -E -q "^(0\.0\.0\.0/1|0/1)[[:space:]]+.*$ts_iface"; then
      echo "  [] Default route successfully overridden via $ts_iface."
    else
      echo "  [!] Failed to verify exit node routing on $ts_iface."
    fi
  else
    echo "[Warning] Could not detect Tailscale utun interface for host routing."
  fi
}

# Helper to nuke leftover network state after teardown (proxy settings, routing
# table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
# problem where stale state kills internet after the gateway stops.
cleanup_network_state() {
  echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."

  # 1. Disable any leftover proxy settings (needs root on many macOS versions)
  for svc in "$ACTIVE_SERVICE" "$EN0_SERVICE"; do
    disable_all_proxies "$svc"
  done
  echo "  Proxies disabled."

```

```
  # 2. Flush DNS cache
  if command -v dscacheutil >> output.log 2>&1; then
    sudo -n dscacheutil -flushcache 2>> output.log || true
  fi
  sudo -n killall -HUP mDNSResponder 2>> output.log || true
  echo "  DNS cache flushed."

  # 3. Flush stale routing table entries
  if command -v route >> output.log 2>&1; then
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
  fi
  remove_warp_bypass_routes
  disable_killswitch
  echo "  Routing table and firewall flushed."

  # 4. Power-cycle Wi-Fi to clear any lingering interface state
  WIFI_PORT=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print $2
    }')
  if [ -n "$WIFI_PORT" ]; then
    sudo -n networksetup -setairportpower "$WIFI_PORT" off 2>> output.log || true
    sleep 2
    sudo -n networksetup -setairportpower "$WIFI_PORT" on 2>> output.log || true
    echo "  Wi-Fi power-cycled (interface $WIFI_PORT)."
    sleep 3
  else
    echo "  Wi-Fi interface not detected; bouncing en0..."
    sudo -n ifconfig en0 down 2>> output.log || true
    sudo -n ifconfig en0 up 2>> output.log || true
  fi

  # Force restart sharing services to prevent AirDrop / AirPlay discovery freezes
  # after network interface/DNS changes.
  echo "  Resetting macOS sharing services (AirDrop/AirPlay)..."
  reset_sharing_services

  echo "Network state cleanup complete."
}

SOCKS_PROXY_PORT=${SOCKS_PROXY_PORT}

# Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
# during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
# `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
# would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.


echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
reset_dns




echo "Checking Gateway Status..."
HIJACK_HOST=$(read_env_var GATEWAY_HIJACK_HOST | tr '[:upper:]' '[:lower:]')

# Decide STOP vs START from persisted intent (marker), not a live Docker probe.
# This makes *disable* instant and robust even when Colima/Docker is down.
#
# CRITICAL (set -e + set -x on bash 3.2): gateway_state ECHOES "running"/"stopped"
# and always returns 0, so we branch on the STRING  never on a function's exit
# code. A helper that `return 1`s trips a spurious `set -e` abort under `set -x`
# (DEBUG_TRACE=true) even when guarded, hence the echo contract. See devref 15.11.8.
if [ "$(gateway_state)" = "running" ]; then
  TOGGLE_PHASE="stop"
  echo -e "\n=============================================="
  echo "Gateway is RUNNING. Stopping it now..."
  echo -e "==============================================\n"

  # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
  if [[ "$HIJACK_HOST" != "false" ]]; then
    disconnect_tailscale_host
    echo ""
  fi

  echo "Stopping Docker containers..."
  # `|| true`: a disable must always complete its teardown even if the Docker
  # daemon / Colima VM is already down (compose down would exit non-zero and,
  # under set -e, abort the STOP). The START path intentionally does NOT guard
  # its `docker compose up`  a failed start SHOULD trigger cleanup.
  run_logged docker compose down --remove-orphans -t 30 || true
```

```
1549
1550    # Only stop Colima if the user explicitly opted in (false by default) to prevent breaking their other Docker
          dev projects
1551    STOP_COLIMA=$(read_env_var STOP_COLIMA_ON_EXIT | tr '[:upper:]' '[:lower:]')
1552    if [[ "$STOP_COLIMA" == "true" ]]; then
1553      echo -e "\nStopping Colima VM to free up host RAM and battery..."
1554      run_with_timeout 30 colima stop >> output.log 2>&1 || echo "Warning: Failed to stop Colima gracefully."
1555    else
1556      echo -e "\nLeaving Colima running. (Set STOP_COLIMA_ON_EXIT=true in .env to change this)"
1557    fi
1558
1559    if [[ "$HIJACK_HOST" != "false" ]]; then
1560      # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
1561      # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
1562      # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
1563      echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
1564      reset_dns
1565
1566      # 3. Stop local DNS proxy if running
1567      stop_local_dns_proxy
1568
1569      # Stop background daemons
1570      stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
1571      stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
1572      stop_warp_watcher
1573      clear_gateway_active_marker
1574
1575      # 4. Nuke leftover network state so internet actually works after teardown
1576      cleanup_network_state
1577    fi
1578
1579    ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
1580    echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
1581  else
1582    TOGGLE_PHASE="start"
1583    START_GATEWAY=true
1584    echo "[$(date '+%Y-%m-%d %H:%M:%S')] Action: STARTUP GATEWAY" >> "$LOG_FILE"
1585    echo -e "\n=============================================="
1586    echo "Gateway is STOPPED. Starting it now..."
1587    echo -e "==============================================\n"
1588
1589    # 1. Prevent host exit-node deadlock during VM / Container startup
1590    if [[ "$HIJACK_HOST" != "false" ]]; then
1591      disconnect_tailscale_host
1592    fi
1593
1594    # 2. Wait for physical DHCP lease to settle (crucial for restarts/wake-up/roaming)
1595    wait_for_dhcp_settle
1596
1597    # 3. Boot Colima VM if it is not already running
1598    echo -e "\nChecking Colima VM status..."
1599    if ! run_with_timeout 15 colima status >> output.log 2>&1; then
1600      # Colima networking is gated on the SAME LAN-P2P signal routing-fix uses
1601      # (.lan_p2p_detected overrides TAILSCALE_ALLOW_LAN_P2P in .env). A LAN-reachable
1602      # VM (--network-address + bridged/shared) needs the socket_vmnet helper and is
1603      # only worthwhile on a trusted home network, so we request it only when P2P is
1604      # allowed; otherwise  AP-isolated / enterprise, the common case  we boot plain
1605      # user-mode networking. Either way host<->container P2P is UNAFFECTED: it rides
1606      # routing-fix's Docker-gateway / .host_ips RETURN rules, not the VM network mode
1607      # (see devref 15.3.5). colima_start_until_ready returns as soon as Docker answers
1608      # (~15s) instead of blocking on colima's ~2min post-provision hang (see 15.8.7);
1609      # that hang is colima-internal and mode-independent (it also occurs in user mode).
1610      allow_p2p=$(read_env_var "TAILSCALE_ALLOW_LAN_P2P")
1611      [ -z "$allow_p2p" ] && allow_p2p="false"
1612      if [ -f "$SCRIPT_DIR/.lan_p2p_detected" ]; then
1613        allow_p2p=$(tr -d '[:space:]' < "$SCRIPT_DIR/.lan_p2p_detected" 2>/dev/null || echo "false")
1614      fi
1615
1616      if [ "$allow_p2p" = "true" ]; then
1617        colima_network_mode="bridged"
1618        if [[ "$OSTYPE" == "darwin"* ]]; then
1619          security_type=$(system_profiler SPAirPortDataType 2>/dev/null | awk '/Current Network Information:/{
      found=1} found && /Security:/{print $NF; exit}')
1620          if echo "$security_type" | grep -qi "Enterprise"; then
1621            echo -e "\n  [!] WPA2-Enterprise Wi-Fi detected! Bridged mode will cause MAC-spoofing
      deauthentication."
1622            echo "      Falling back to '--network-mode shared' for Colima."
1623            colima_network_mode="shared"
1624          fi
1625        fi
```

```
      echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P on: --network-address $colima_network_mode)..."
      colima_start_until_ready --memory 0.6 --vm-type vz --network-address --network-mode "$colima_network_mode" \
        || echo "  [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry)."
    else
      colima_network_mode="user"
      echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P off: fast user-mode networking)..."
      colima_start_until_ready --memory 0.6 --vm-type vz \
        || echo "  [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry)."
    fi
    echo "$colima_network_mode" > /tmp/nullexit_colima_mode.txt
  else
    echo "Colima is already running."
  fi

  # 3b. Auto-detect and fix Docker Subnet Collisions (e.g. if the user travels to a 172.17.x.x Wi-Fi)
  if [ -f "scripts/fix-docker-bridge-collision.sh" ]; then
    if bash scripts/fix-docker-bridge-collision.sh --dry | grep -q "collision detected"; then
      echo -e "\n[!] WARNING: Docker Subnet Collision Detected with your current Wi-Fi!"
      echo -e "    Auto-fixing Colima config to prevent internet jitter..."
      bash scripts/fix-docker-bridge-collision.sh --skip-toggle || true
    fi
  fi

  # 3c. Configure swap file inside the VM to prevent OOM on low-memory limits
  # We also set vm.swappiness=10 so the Linux kernel strictly prefers physical RAM
  # and avoids unnecessarily wearing out the SSD with proactive background swapping.
  if ! run_with_timeout 15 colima ssh -- grep -q 'swapfile' /proc/swaps >> output.log 2>&1; then
    echo "Configuring 400MB swap file inside the VM to prevent OOM..."
    run_with_timeout 30 colima ssh -- sudo sh -c "if [ ! -f /swapfile ]; then dd if=/dev/zero of=/swapfile bs=1M count=400 status=none && mkswap /swapfile; fi && swapon /swapfile && sysctl vm.swappiness=10" >> output.log 2>&1 || echo "Warning: Failed to enable swap file inside the VM."
  fi

  # 4. Clean up corrupted AdGuardHome configurations
  if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
    rm -f "adguard/conf/AdGuardHome.yaml"
    echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
  fi

  # 4b. Auto-heal stale Docker socket from hard reboots
  if ! run_with_timeout 15 docker ps >/dev/null 2>&1; then
    echo "Docker daemon is unresponsive (likely a stale socket from a crash). Auto-healing Colima..."
    run_with_timeout 120 colima restart >> output.log 2>&1
  fi

  # 4c. Inject TCP MSS clamp from .env into post-rules.txt before starting
  if [ -f .env ] && grep -q "^GATEWAY_MSS=" .env; then
    GATEWAY_MSS=$(read_env_var GATEWAY_MSS)
    if [[ "$GATEWAY_MSS" =~ ^[0-9]+$ ]]; then
      # macOS sed syntax
      sed -i '' "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null || \
      # Linux sed fallback
      sed -i "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null
    fi
  fi

  # 5. Start compose services
  write_host_ips

  # Procedurally generated ports have already been exported at the top of the script.
  echo "  [Security] Procedurally randomized proxy ports generated to evade fingerprinting."

  # --- HONEY-PORT TRIPWIRE ---
  # Generate a random trap port. If any malware does a sequential port scan on localhost,
  # it will inevitably hit this port. When it does, we instantly fire a macOS alert.
  HONEY_PORT=$(get_free_port)
  (
    if nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application just pinged the local Honey-Port ($HONEY_PORT)!" >> "$LOG_FILE"
      osascript -e 'display notification "An unknown application is aggressively scanning your local ports. Malware port-scan suspected." with title "nullexit OPSEC Alert" sound name "Basso"' 2>/dev/null || true
    fi
  ) &
  disown %1 2>/dev/null || true
  echo "  [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."
```

```
1698
1699  echo -e "\nStarting Docker containers..."
1700
1701  # --- Ad-block rule compilation: hash-gated one-off, off the critical path (15.8.8) ---
1702  # The rule-compiler (a `compile`-profiled container, the SOLE writer of adguard/work)
1703  # re-dedupes ~340k rules inside the 600MB VM (~28s)  far too slow to run every toggle.
1704  # It is gated on a hash of the lists YOU control (black_list.txt + white_list.txt):
1705  # recompile only when they change, when compiled_rules.txt / ip_blocklist.ipset are
1706  # missing, or when they exceed RULES_MAX_AGE_DAYS (default 7, so the remote blocklists
1707  # still refresh occasionally). Otherwise AdGuard/routing-fix boot on the existing
1708  # artifacts and we skip the ~28s  they no longer `depends_on` the rule-compiler
1709  # container (see docker-compose.yml). A compile failure is non-fatal: we keep the
1710  # previous compiled rules rather than bricking the tunnel (ad-block != the kill-switch).
1711  RULES_HASH_FILE="$SCRIPT_DIR/.rules_hash"
1712  rules_hash=$(compute_rules_hash)
1713  stored_rules_hash=""
1714  [ -f "$RULES_HASH_FILE" ] && stored_rules_hash=$(cat "$RULES_HASH_FILE" 2>/dev/null)
1715  rules_max_age_days=$(read_env_var RULES_MAX_AGE_DAYS)
1716  [ -z "$rules_max_age_days" ] && rules_max_age_days=7
1717  need_compile=false
1718  if [ ! -f adguard/work/userfilters/compiled_rules.txt ] || [ ! -f adguard/work/userfilters/ip_blocklist.ipset
         ]; then
1719    need_compile=true
1720  elif [ -z "$rules_hash" ] || [ "$rules_hash" != "$stored_rules_hash" ]; then
1721    need_compile=true
1722  elif [ -n "$(find adguard/work/userfilters/compiled_rules.txt -mtime +"$rules_max_age_days" 2>/dev/null)" ];
       then
1723    need_compile=true
1724  fi
1725  if [ "$need_compile" = "true" ]; then
1726    echo "  Ad-block lists changed (or artifacts missing/stale)  compiling rules (one-off, ~28s)..."
1727    if run_logged docker compose run --build --rm rule-compiler >> output.log 2>&1; then
1728      if [ -n "$rules_hash" ]; then echo "$rules_hash" > "$RULES_HASH_FILE"; fi
1729    else
1730      echo "  [!] Rule compile failed  continuing with existing compiled ad-block rules."
1731    fi
1732  else
1733    echo "  Ad-block lists unchanged  reusing compiled rules (skipping the ~28s recompile)."
1734  fi
1735
1736  # --- Gate `--build` for the long-running images (routing-fix, tor) ---
1737  # BuildKit re-checking these on the 600MB VM costs ~30s even when every layer is
1738  # CACHED (15.8.8). Only pass --build when the hashed inputs change or an image is
1739  # missing; otherwise plain `up -d` reuses cached images. `up` starts everything
1740  # except the `compile`-profiled rule-compiler.
1741  BUILD_HASH_FILE="$SCRIPT_DIR/.build_hash"
1742  build_inputs_hash=$(compute_build_inputs_hash)
1743  stored_build_hash=""
1744  [ -f "$BUILD_HASH_FILE" ] && stored_build_hash=$(cat "$BUILD_HASH_FILE" 2>/dev/null)
1745  local_images_present=true
1746  for _img in nullexit-routing-fix:v1.0.0 nullexit-tor:v1.0.0; do
1747    docker image inspect "$_img" >> output.log 2>&1 || local_images_present=false
1748  done
1749  if [ -n "$build_inputs_hash" ] && [ "$build_inputs_hash" = "$stored_build_hash" ] && [ "$local_images_present
       " = "true" ]; then
1750    echo "  Build inputs unchanged  skipping image build (reusing cached images)."
1751    run_logged docker compose up -d
1752  else
1753    echo "  Build inputs changed or a local image is missing  building images..."
1754    run_logged docker compose up -d --build
1755    # Record the hash only after a successful build (set -e aborts above on failure).
1756    if [ -n "$build_inputs_hash" ]; then echo "$build_inputs_hash" > "$BUILD_HASH_FILE"; fi
1757  fi
1758
1759  RULE_COUNT=$(grep "! Total Block Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -F': '
       '{print $2}' || echo "0")
1760  NATIVE_COUNT=$(grep "! Native AdGuard Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -
       F': ' '{print $2}' || echo "0")
1761
1762  if [ "$RULE_COUNT" != "0" ] && [ "$RULE_COUNT" != "" ]; then
1763    if [ "$NATIVE_COUNT" != "0" ] && [ "$NATIVE_COUNT" != "" ]; then
1764      TOTAL_COUNT=$((RULE_COUNT + NATIVE_COUNT))
1765      # Format with commas for readability (e.g. 441,578)
1766      if command -v printf &> /dev/null; then
1767        FORMATTED_RULE=$(printf "%'d" "$RULE_COUNT")
1768        FORMATTED_TOTAL=$(printf "%'d" "$TOTAL_COUNT")
1769        echo "  DNS: Compiled $FORMATTED_RULE unique custom rules (Total active in AdGuard: ~$FORMATTED_TOTAL
       rules)."
1770      else
1771        echo "  DNS: Compiled $RULE_COUNT unique custom rules (Total active in AdGuard: ~$TOTAL_COUNT rules)."
1772      fi
```

```
    else
      echo "  DNS: Compiled and loaded $RULE_COUNT optimized rules."
    fi
  fi

  # Report IP blocklist compilation results
  IP_COUNT=$(grep "^# Entries:" adguard/work/userfilters/ip_blocklist.ipset 2>/dev/null | awk -F': ' '{print $2
     }' || echo "0")
  if [ "$IP_COUNT" != "0" ] && [ "$IP_COUNT" != "" ]; then
    if command -v printf &> /dev/null; then
      FORMATTED_IP=$(printf "%'d" "$IP_COUNT")
      echo "  IP:  Loaded $FORMATTED_IP threat intelligence IPs/CIDRs into kernel firewall."
    else
      echo "  IP:  Loaded $IP_COUNT threat intelligence IPs/CIDRs into kernel firewall."
    fi
  fi

  # 6. Wait for the gateway container's Tailscale connection to be ready
  BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
  USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')

  echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
  echo "  (This can take 30-60s  Tailscale goes through: NoState  Starting  Running  Online)"
  connected=false
  consecutive_failures=0
  for i in {1..60}; do
    # Run tailscale status and capture its output so the user can see what's happening.
    # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
    ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)

    if echo "$ts_output" | grep -q "offers exit node"; then
      connected=true
      break
    fi

    # Track consecutive failures to detect a stuck state.
    # If we see 40+ consecutive 'NoState' responses, give up early
    # the daemon is alive but can't connect to the coordination server.
    if echo "$ts_output" | grep -q "NoState"; then
      consecutive_failures=$((consecutive_failures + 1))
      if [ "$consecutive_failures" -ge 40 ]; then
        echo ""
        echo "  [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
        echo "    Tailscale daemon is running but can't reach the coordination server."
        break
      fi
    else
      consecutive_failures=0
    fi

    # Print a condensed status every 10 seconds so the user can see progress
    if [ $((i % 10)) -eq 0 ]; then
      ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
      echo "  [attempt $i/60] $ts_short"
    elif [ $((i % 5)) -eq 0 ]; then
      echo -n "[$i]"
    else
      echo -n "."
    fi
    sleep 1
  done
  echo ""

  if [ "$connected" = "true" ]; then
    echo "Gateway container is online on the Tailscale mesh."
  elif [ "$BYPASS_PING" = "true" ]; then
    echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
  else
    echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
    echo "  The container was seen in the following states during the wait:" >&2
    echo "$ts_output" | sed 's/^/    /' >&2
    echo "" >&2
    echo "  This could mean:" >&2
    echo "    1. Tailscale auth key has expired  check TS_AUTHKEY in .env" >&2
    echo "    2. Network issue inside the container  check: docker compose logs tailscale" >&2
    echo "    3. gluetun VPN tunnel not connecting  check: docker compose logs warp | tail -20" >&2
    echo "" >&2
    echo "  Diagnose manually:" >&2
    echo "    docker compose exec -T tailscale tailscale status" >&2
    echo "    docker compose exec -T tailscale tailscale netcheck" >&2
    cleanup_handler ERR $LINENO
```

```
    exit 1
  fi

  # 6b. Resolve the gateway's Tailscale IP.
  # Dynamically query the container for the IP (handles IP changes after re-auth)
  # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
  #
  # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
  # returns (\r) into captured output. If $TS_IP ends with \r, then
  # `networksetup -setdnsservers` silently rejects the malformed IP and
  # DNS stays at 1.1.1.1  the primary bug this project was hitting.
  # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
  # the first field of the first line.
  TS_IP=""
  if [ "$connected" = "true" ]; then
    TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
     awk 'NR==1{print $1; exit}' || true)
    if [ -n "$TS_IP" ]; then
      echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
      echo "$TS_IP" > .gateway_ip
    else
      echo "    [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
      TS_IP=$(read_adguard_ip || true)
    fi
  else
    TS_IP=$(read_adguard_ip || true)
  fi

  # 7. Connect host Mac to Tailscale mesh.
  # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
  # system LaunchDaemon via 'brew services start tailscale'), the previous five-
  # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
  # no menu bar item to click, and no Accessibility permission required.
  #
  # CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
  # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
  # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
  # are impossible without the host on the mesh).
  #
  # We connect in two phases:
  #   Phase A  Join mesh WITHOUT exit node (so pre-flight checks can run)
  #   Phase B  Set exit node only after pre-flight checks confirm the gateway works
  HOST_ON_MESH=false
  SKIP_EXIT_NODE=false
  if [ "$HIJACK_HOST" = "false" ]; then
    echo -e "\n[Info] HIJACK_HOST is false (Headless Mode). Host networking will remain untouched."
    SKIP_EXIT_NODE=true
  elif [ -n "$TS_BIN" ]; then
    echo -n "Verifying tailscaled is reachable"
    daemon_ready=false
    status_exit=0
    for i in {1..10}; do
      # Test if the daemon responds to status check.
      # If status returns non-zero, check if it was a normal "stopped/disconnected" state
      # (exit code 1 is normal for disconnected). If the command exited quickly (not killed by timeout),
      # and the error is just "Tailscale is stopped", then the daemon is alive and ready.
      # But if it timed out (exit code 143) or is completely unresponsive, it's wedged.
      status_exit=0
      if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
        daemon_ready=true
        break
      else
        status_exit=$?
        if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
          daemon_ready=true
          break
        fi
      fi
      echo -n "."
      sleep 1
    done
    echo ""

    # If the daemon was unresponsive or socket check failed, attempt auto-restart
    if [ "$daemon_ready" != "true" ]; then
      if [ "$status_exit" -eq 143 ]; then
        warn "tailscaled daemon is wedged/unresponsive. Restarting it..."
      else
        warn "tailscaled daemon is not running. Attempting to start it..."
      fi
      restart_tailscaled_daemon
```

```
      # Re-verify
      if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
        daemon_ready=true
      fi
    fi

    if [ "$daemon_ready" = "true" ]; then
      echo "tailscaled is responsive (running as a system service)."

      #  Phase A: Join mesh without exit node
      echo "Connecting host to Tailscale mesh (no exit node yet)..."
      if $TS_BIN up; then
        $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node= || true
        HOST_ON_MESH=true
        echo "Host is on Tailscale mesh."
      else
        warn "tailscale up failed even though the daemon is running."
        echo "  This usually means the host hasn't authenticated with your tailnet yet."
        echo "  Run this command in a terminal to authenticate:"
        echo ""
        echo "      sudo tailscale up"
        echo ""
        echo "  A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
      fi
    else
      fail "tailscaled could NOT be started automatically."
      echo "  Run this in a terminal to start and authenticate Tailscale:"
      echo ""
      echo "      sudo brew services start tailscale"
      echo "      sudo tailscale up"
      echo ""
      echo "  Then re-run toggle.sh."
    fi
  else
    echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
  fi

  #  Phase B: Pre-flight checks + enable exit node only if safe
  if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
    echo -e "\n"
    echo "Pre-flight connectivity check"
    echo ""

    # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
    # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
    # ("direct connection not established") even though the pong was received.
    # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
    echo -n "  [1/3] Gateway reachable via Tailscale... "
    ping_ok=false
    # The host tailscaled needs a few seconds to propagate the new network map and
    # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
    for attempt in {1..45}; do
      if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
        ping_ok=true
        break
      fi
      sleep 1
    done
    if [ "$ping_ok" = "true" ]; then
      echo "PASS"
    else
      echo "FAIL"
      SKIP_EXIT_NODE=true
    fi

    # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
    # Colima's SSH tunnel only forwards TCP (not UDP), so we use +tcp.
    echo -n "  [2/3] AdGuard DNS via localhost... "
    if command -v dig &>/dev/null; then
      if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 &>/dev/null; then
        echo "PASS"
      else
        echo "FAIL"
        SKIP_EXIT_NODE=true
      fi
    elif command -v nslookup &>/dev/null; then
      if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 &>/dev/null; then
        echo "PASS"
      else
        echo "FAIL"
```

```
        SKIP_EXIT_NODE=true
      fi
    else
      echo "SKIP (dig/nslookup not available)"
    fi

    # Check 3: Does the WARP container have external internet?
    echo -n "  [3/3] WARP container internet... "
    tmp_err=$(mktemp)
    if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace >/dev/null 2>"$tmp_err"; then
      echo "PASS"
    else
      echo "FAIL"
      if [ -s "$tmp_err" ]; then
        err_msg=$(cat "$tmp_err" | tr -d '\r')
        echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Check 3] WARP container check failed: $err_msg" >> output.log
      fi
      SKIP_EXIT_NODE=true
    fi
    rm -f "$tmp_err"

    # Act based on check results
    if [ "$SKIP_EXIT_NODE" != "true" ]; then
      echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
      $TS_BIN up || true
      if $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node="$TS_IP" --exit-node-allow-lan-access=true; then
        echo "Exit node enabled."
        add_warp_bypass_routes
        setup_exit_node_routing
        enable_killswitch
      else
        echo "[Warning] Failed to set exit node."
        SKIP_EXIT_NODE=true
      fi
    fi

    if [ "$SKIP_EXIT_NODE" = "true" ]; then
      warn "Pre-flight checks failed  EXIT NODE + DNS HIJACK SKIPPED."
      echo "  Host stays on Tailscale mesh but your internet routes through your normal connection."
      echo "  Troubleshoot with:"
      echo "    docker compose logs warp | tail -20"
      echo "    docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"
      echo ""
      echo "  Falling back to SOCKS5 proxy for traffic routing..."
    fi
  elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
    echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
    SKIP_EXIT_NODE=true
  fi

  # 8. Apply DNS Hijacking.
  # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
  # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
  # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
  if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
    echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
    echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
    if force_dns_to_gateway "$TS_IP"; then
      echo "DNS hijack verified: single server = $TS_IP."
    else
      echo "[Warning] DNS hijack could not be verified."
      echo "  If DNS is broken, run manually:"
      echo "    networksetup -setdnsservers \"$ACTIVE_SERVICE\" $TS_IP"
      echo "    networksetup -setsearchdomains \"$ACTIVE_SERVICE\" ts.net"
    fi
  elif [ -n "$TS_IP" ] && [ "$HIJACK_HOST" != "false" ]; then
    echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
    echo "  Trying local DNS proxy for ad-blocking..."
    if start_local_dns_proxy; then
      echo "  Ad-blocking active via local DNS proxy through AdGuard."
    else
      echo "  Local DNS proxy unavailable. DNS stays at 1.1.1.1."
      echo "  AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
    fi

    # Enable SOCKS5 proxy for traffic routing through WARP
    # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
    echo -e "\n  Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
    echo "  (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
```

```
    enable_socks_proxy "$ACTIVE_SERVICE"

    # Verify the SOCKS5 proxy works through WARP
    echo -n "  Verifying SOCKS5 proxy routes through WARP... "
    if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn-cgi/
     trace 2>/dev/null | grep -q "warp=on"; then
      echo "ok (warp=on)."
      echo "  All TCP traffic now encrypted through Cloudflare WARP tunnel."
    else
      echo "FAILED (proxy not routing through WARP)."
      echo "  Disabling SOCKS5 proxy  internet stays on direct connection."
      disable_socks_proxy
      echo "  SOCKS proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
    fi
  else
    echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
  fi

  # 9. Verify host connectivity
  if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
    if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
      echo "Host is online on the Tailscale mesh."
    else
      echo "[Warning] Host is not responding on Tailscale."
    fi
  fi

  ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
  echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."

  # Prevent system from going to sleep while gateway is running
  start_sleep_prevention

  if [[ "$HIJACK_HOST" != "false" ]]; then
    if [ -n "$TS_IP" ] && [ "$TS_IP" != "1.1.1.1" ]; then
      start_dns_watcher "$TS_IP"
    fi

    # Start WARP liveness monitor (logs to output.log on warp flip)
    start_warp_watcher
  fi

  # Reset sharing services after network configuration to prevent AirDrop freezes
  echo "Resetting macOS sharing services (AirDrop/AirPlay)..."
  reset_sharing_services

  # Tell external watchers (post-wake / network-change) the gateway is up.
  write_gateway_active_marker

  #  Local Network Surveillance Check
  echo -e "\n"
  echo "Local Network Surveillance Check"
  echo ""
  # Count populated ARP cache entries to detect network scanning or high congestion
  # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
  ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
  if [ "$ARP_COUNT" -gt 15 ]; then
    warn "High ARP activity detected ($ARP_COUNT devices in cache)."
    warn "This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)."
    echo -e "  [] Your traffic remains fully encrypted and invisible.\n"
  else
    echo -e "  [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
  fi
fi

SUCCESS_RUN=true

# Final DNS state summary
if [ -n "$ACTIVE_SERVICE" ]; then
  FINAL_DNS=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log || true)
  echo -e "\n"
  echo -e "DNS STATE: $FINAL_DNS"
  echo -e ""
fi

if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
  if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
    echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
    echo "  The gateway gracefully fell back to yesterday's cached blocklist."
    echo "  (Run 'docker logs rule-compiler' later to investigate which link died.)"
  fi
```

```
fi

rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
if [ -t 0 ]; then
  read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
  if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
    echo "Reversing gateway state..."
    exec bash "$0"
  fi
fi

echo -e "\nYou can close this terminal window now."
```

7 recover.sh

```
#!/bin/bash
# recover.sh  nullexit KILL-SWITCH recovery script (macOS + Linux)
# Fixes internet when toggle.sh leaves you stranded.
#
# This is a NUCLEAR option: it undoes EVERYTHING toggle.sh does by:
#   1. Disconnecting host Tailscale from the mesh (exit-node off)
#   2. Resetting DNS to default on ALL network services
#   3. Disabling rogue proxy settings (SOCKS, web, secure web)
#   4. Flushing the DNS cache
#   5. Flushing stale routing table entries
#   6. Stopping gateway Docker containers
#   7. Power-cycling Wi-Fi
#   8. Verifying internet is working
#
# One file, both OSes: the dispatch below runs the self-contained Linux
# recovery and exits; on macOS it falls through to the richer dual-mode
# implementation (default teardown / --post-wake refresh).
#
# Usage:  double-click Recover-Gateway.applescript
#         OR  ./recover.sh
#         OR  ./recover.sh --post-wake    (macOS only  lighter refresh, keeps gateway live)

set -e

#  Resolve script directory (used for path-relative refs)
# recover.sh lives at the repo root; SCRIPT_DIR lets us launch toggle.sh
# (and reference any other repo-root file) from the same directory no
# matter where the user invoked recover.sh from.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
cd "$SCRIPT_DIR" || exit 1

#
# LINUX RECOVERY  self-contained; runs the full teardown then exits before the
# macOS body below. Native Linux tooling (resolvectl / nmcli / ip) throughout.
#
if [[ "$OSTYPE" != "darwin"* ]]; then
  # Source common formatting and helper functions
  source "$SCRIPT_DIR/scripts/common.sh"

  # Always add Homebrew path
  export PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:$PATH"

  echo ""
  echo -e "${RED}${BOLD}${NC}"
  echo -e "${RED}${BOLD}     Nullexit Gateway  RECOVERY MODE     ${NC}"
  echo -e "${RED}${BOLD}${NC}"
  echo ""
  echo "This script will tear down the gateway and restore normal internet."
  echo ""

  #  Cache sudo credentials upfront
  # Prevents the script from hanging halfway through when it needs to run
  # privileged commands (route flush, Wi-Fi power cycle, etc.).
  echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
  sudo -v
  (
    while true; do
      sudo -n true 2>/dev/null
      sleep 60
    done
  ) &
```

```
  SUDO_KEEPER_PID=$!
  trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT

  #  1. Disconnect host Tailscale
  step "Disconnecting host Tailscale from mesh"
  if command -v tailscale >> output.log 2>&1; then
    # Check if actually connected first (with timeout  tailscaled could be wedged)
    if run_with_timeout 5 tailscale status >> output.log 2>&1; then
      # Also explicitly reset any exit-node so it doesn't linger.
      if run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= >> output.log
   2>&1; then
        ok "Exit-node preference cleared"
      else
        warn "tailscale up --reset didn't respond (tailscaled may be wedged)"
      fi
      ts_args=""
      if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
      if run_with_timeout 10 tailscale down $ts_args >> output.log 2>&1; then
        ok "Tailscale disconnected"
      else
        warn "tailscale down didn't respond"
      fi
    else
      warn "tailscaled not reachable  skipping (timeout after 5s)"
    fi
  else
    warn "tailscale CLI not found  skipping"
  fi

  #  2. Reset DNS to default on ALL network services
  step "Resetting DNS to default on active interface"

  ACTIVE_SERVICE=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
  if [ -n "$ACTIVE_SERVICE" ]; then
    sudo resolvectl revert "$ACTIVE_SERVICE" 2>> output.log || true
    ok "DNS reset on $ACTIVE_SERVICE"
  else
    warn "Could not determine active network interface"
  fi

  #  3. Disable rogue proxy settings
  step "Disabling proxy settings (SOCKS, web, secure web)"
  echo "  (Linux global SOCKS proxy configuration skipped.)"
  ok "Proxy settings cleared"


  #  4. Flush DNS cache
  step "Flushing DNS cache"
  if command -v resolvectl >> output.log 2>&1; then
    sudo resolvectl flush-caches 2>> output.log || true
    ok "DNS cache flushed"
  else
    warn "resolvectl not found"
  fi

  #  5. Clear stale routing table
  step "Flushing stale routing table entries"
  sudo ip route flush cache >> output.log 2>&1 || true
  ok "Routing table flushed"

  #  6. Stop gateway Docker containers
  step "Stopping gateway Docker containers"
  if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
    if [ -f "docker-compose.yml" ]; then
      docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
    else
      warn "docker-compose.yml not found in current directory"
    fi
  else
    warn "Docker not available  skipping"
  fi

  #  6b. Stopping sleep prevention (systemd-inhibit)
  step "Stopping sleep prevention"
  PID_FILE="/tmp/nullexit-caffeinate.pid"
  if [ -f "$PID_FILE" ]; then
    INHIBIT_PID=$(cat "$PID_FILE")
    if [ -n "$INHIBIT_PID" ] && kill -0 "$INHIBIT_PID" 2>/dev/null; then
      kill "$INHIBIT_PID" 2>/dev/null || true
      ok "Sleep prevention stopped (PID $INHIBIT_PID)"
    else
```

```
      warn "Sleep prevention process not running, cleaning up stale PID file"
    fi
    rm -f "$PID_FILE"
  else
    ok "Sleep prevention was not active"
  fi

  #  7. Power-cycle Wi-Fi
  step "Power-cycling Wi-Fi"
  if command -v nmcli &>/dev/null; then
    sudo nmcli radio wifi off 2>> output.log || true
    sleep 2
    sudo nmcli radio wifi on 2>> output.log || true
    ok "Wi-Fi bounced via nmcli"
    sleep 3
  else
    warn "nmcli not found. Falling back to ip link..."
    WIFI_IFACE=$(ip link | grep -E '^[0-9]+: (wl[a-zA-Z0-9]+|wlan[0-9]+):' | awk -F': ' '{print $2}')
    if [ -n "$WIFI_IFACE" ]; then
      sudo ip link set "$WIFI_IFACE" down 2>> output.log || true
      sleep 2
      sudo ip link set "$WIFI_IFACE" up 2>> output.log || true
      ok "Wi-Fi bounced via ip link (interface $WIFI_IFACE)"
      sleep 3
    else
      warn "No Wi-Fi interface detected."
    fi
  fi

  #  8. Verify internet connectivity
  verify_internet_connectivity

  #  Summary
  echo ""
  echo -e "${BOLD}${NC}"
  echo -e "${BOLD}RECOVERY SUMMARY${NC}"
  echo -e "${BOLD}${NC}"

  # Show final DNS state
  FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr
     '\n' ' ' || echo "N/A")
  echo "  DNS ($ACTIVE_SERVICE):  $FINAL_DNS"

  echo ""

  if [ "$INTERNET_OK" = "true" ]; then
    echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
  else
    echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
    echo "    This might mean:"
    echo "    - Your Wi-Fi is disconnected from the router"
    echo "    - Tailscale is still interfering (try restarting tailscaled:"
    echo "        systemctl restart tailscaled)"
    echo "    - You need to re-connect to Wi-Fi in the menu bar"
    echo "    - Try toggling Wi-Fi off/on in the menu bar"
    echo ""
    echo "    Or try manually:"
    echo "      resolvectl revert $ACTIVE_SERVICE"
    extra_args=""
    if is_kill_switch_enabled; then extra_args="--accept-risk=lose-ssh"; fi
    echo "      tailscale down $extra_args"
    echo "      sudo ip route flush cache"
    echo "      systemctl restart tailscaled"
  fi

  echo ""
  echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
  echo ""

  exit 0
fi

#
# macOS RECOVERY (dual-mode: default teardown / --post-wake refresh)
#

# Enforce Cryptographic Script Integrity
if [ -f "scripts/crypto.sh" ]; then
  if ! bash scripts/crypto.sh --verify; then
    exit 1
  fi
```

```
fi

#  Argument parsing
# --post-wake: lightweight refresh after sleep/wake or Wi-Fi roam.
#              Keeps the gateway live while HUPing DNS caches, refreshing the
#              Tailscale exit-node leak, re-hijacking host DNS, and force-
#              recreating the warp container if its gluetun healthcheck failed.
#              Does NOT tear down Docker, Tailscale, sleep prevention, or Wi-Fi.
#              Called from scripts/watcher.sh on every wake / network-change event
#              while /tmp/nullexit-gateway-active.marker is present (i.e. the
#              gateway is currently up). See devref.md 10.29 for the why.
POST_WAKE=false
AUTO_MODE=false
for arg in "$@"; do
  if [ "$arg" = "--post-wake" ]; then
    POST_WAKE=true
  elif [ "$arg" = "--auto" ] || [ "$arg" = "--non-interactive" ]; then
    AUTO_MODE=true
  fi
done

rm -f "$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker"

# Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
LOCK_FILE="/tmp/nullexit-toggle.lock"
if [ -f "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
      if [ "$POST_WAKE" = "true" ]; then
        echo "[$(date -u +%FT%TZ)] POST-WAKE skipped: another lifecycle script (PID $LOCK_PID) is running." >>
    output.log
        exit 0
      else
        echo "Error: Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running." >&2
        exit 1
      fi
    fi
  fi
fi
echo "$$" > "$LOCK_FILE"

# Clear the active-state marker in nuclear recovery mode since the gateway is being
# completely torn down. This prevents scripts/watcher.sh from triggering post-wake
# recoveries in response to network changes caused by the teardown itself.
if [ "$POST_WAKE" = "false" ]; then
  rm -f "/tmp/nullexit-gateway-active.marker"
fi

# WARP Watcher inhibit marker: written immediately in --post-wake mode so the
# background WARP Watcher (in toggle.sh) skips failure counting while we are
# intentionally tearing down / recreating the warp container. Without this, the
# watcher accumulates 6 consecutive warp=off readings during force-recreate (~35s)
# and fires nuclear recover.sh, killing a gateway that would have been healthy.
WARP_INHIBIT_FILE="/tmp/nullexit-warp-inhibit.marker"
if [ "$POST_WAKE" = "true" ]; then
  touch "$WARP_INHIBIT_FILE"
  echo "[$(date -u +%FT%TZ)] POST-WAKE started  WARP Watcher inhibited to prevent spurious shutdown during warp
      recreate." >> output.log
fi

# Ensure the inhibit marker and lock file are always removed when the script exits (success or error)
cleanup_recover() {
  rm -f "$WARP_INHIBIT_FILE"
  if [ -f "$LOCK_FILE" ] && [ "$(cat "$LOCK_FILE" 2>/dev/null)" = "$$" ]; then
    rm -f "$LOCK_FILE"
  fi
}
trap cleanup_recover EXIT

# Source common formatting and helper functions
source "$SCRIPT_DIR/scripts/common.sh"

# Always add Homebrew path
setup_standard_path

echo ""
if [ "$POST_WAKE" = "true" ]; then
  echo -e "${YELLOW}${BOLD}${NC}"
  echo -e "${YELLOW}${BOLD}  Nullexit  POST-WAKE / POST-ROAM LIGHT ${NC}"
  echo -e "${YELLOW}${BOLD}  (keeps the gateway live, refreshes)    ${NC}"
```

```
  echo -e "${YELLOW}${BOLD}${NC}"
  echo ""
  echo "Re-hijacking DNS, refreshing Tailscale exit-node, force-recreating"
  echo "warp if unhealthy, and resetting sharing services. The gateway stays"
  echo "up. docker compose / Wi-Fi / caffeinate are NOT touched."
  echo ""
else
  echo -e "${RED}${BOLD}${NC}"
  echo -e "${RED}${BOLD}     Nullexit Gateway  RECOVERY MODE     ${NC}"
  echo -e "${RED}${BOLD}${NC}"
  echo ""
  echo "This script will tear down the gateway and restore normal internet."
  echo ""
fi

#  Cache sudo credentials upfront (default mode ONLY)
# Skip in --post-wake: the LaunchAgent-launched watcher has no TTY, so `sudo -v`
# would block forever waiting for a password. All post-wake commands use `sudo -n`
# which is non-blocking and relies on /etc/sudoers.d/nullexit NOPASSWD entries.
SUDO_KEEPER_PID=""
if [ "$POST_WAKE" = "false" ] && [ "$AUTO_MODE" = "false" ] && [ -t 0 ]; then
  echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
  sudo -v
  (
    while true; do
      sudo -n true 2>/dev/null
      sleep 60
    done
  ) &
  SUDO_KEEPER_PID=$!
  trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
fi

#  Resolve physical default interface and gateway upfront
PHYSICAL_IFACE=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
    getline; print $2; exit}')
[ -z "$PHYSICAL_IFACE" ] && PHYSICAL_IFACE="en0"

if [ "$POST_WAKE" = "true" ]; then
  wait_for_dhcp_settle
else
  # Quick resolution in non-post-wake mode (non-blocking)
  PHYSICAL_GW=$(ipconfig getpacket "$PHYSICAL_IFACE" 2>> output.log | awk -F'[{}]' '/router /{print $2}')
fi

#  1. Disconnect host Tailscale (default) / Refresh exit-node (post-wake)

if [ "$POST_WAKE" = "true" ]; then
  step "Refreshing host Tailscale exit-node routing (post-wake)"
  if command -v tailscale >> output.log 2>&1; then
    # Re-issue the SAME exit-node up command toggle.sh START uses. This refreshes
    # the DERP relay mapping and re-asserts the exit-node preference without
    # dropping the host's mesh connection (which `tailscale down` would).
    TS_IP=""
    for src in .gateway_ip; do
      [ -f "$src" ] || continue
      TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
      [ -n "$TS_IP" ] && break
    done

    if [ -n "$TS_IP" ]; then
      step "Re-establishing exit node $TS_IP via 'tailscale set --exit-node'"
      # `tailscale set` only mutates the named preference (exit-node here).
      # Crucially it does NOT call --reset, which would re-apply ALL flags and
      # trigger a sub-second route-table re-evaluation cycle each post-wake.
      # On already-connected nodes, the lighter `set` keeps the existing tunnel
      # alive while only changing the exit-node preference.
      if run_with_timeout 15 tailscale set --exit-node="$TS_IP" \
                           --exit-node-allow-lan-access=true \
                           >> output.log 2>&1; then
        ok "Exit-node $TS_IP re-asserted via 'tailscale set'"
      else
        warn "tailscale set --exit-node=$TS_IP didn't respond; falling back to mesh-only"
        run_with_timeout 10 tailscale set --exit-node= \
                         >> output.log 2>&1 || true
      fi
    else
      warn ".gateway_ip missing  cannot re-assert exit node"
      run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= \
                       >> output.log 2>&1 || true
    fi
```

```
  else
    warn "tailscale CLI not found"
  fi
else
  step "Disconnecting host Tailscale from mesh"
  disconnect_tailscale_host "tailscale"
fi

#  2. Reset DNS to default on ALL network services (default) / Re-hijack gateway DNS (post-wake)
if [ "$POST_WAKE" = "true" ]; then
  step "Re-hijacking host DNS to gateway IP (post-wake / post-roam)"
  TS_IP=""
  for src in .gateway_ip; do
    [ -f "$src" ] || continue
    TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
    [ -n "$TS_IP" ] && break
  done

  if [ -n "$TS_IP" ]; then
    # Detect the active network service (using helper in common.sh)
    ACTIVE_SVC=$(get_active_service)
    EN0_SVC=$(get_en0_service)

    networksetup -setsearchdomains "$ACTIVE_SVC" "ts.net" 2>> output.log || true
    networksetup -setdnsservers "$ACTIVE_SVC" "$TS_IP" 2>> output.log || true
    networksetup -setsearchdomains "$EN0_SVC" "ts.net" 2>> output.log || true
    networksetup -setdnsservers "$EN0_SVC" "$TS_IP" 2>> output.log || true

    HITS=$(networksetup -getdnsservers "$ACTIVE_SVC" 2>> output.log | grep -Fx "$TS_IP" || true)
    if [ -n "$HITS" ]; then
      ok "DNS re-hijacked to $TS_IP on services: $ACTIVE_SVC, $EN0_SVC"
    else
      warn "networksetup accepted but DNS read-back didn't include $TS_IP"
    fi
  else
    warn ".gateway_ip missing  cannot re-hijack DNS (user must run toggle.sh START)"
  fi
else
  step "Resetting DNS to default on all network services (empty = DHCP)"

  # Get ALL network services (handles spaces in names, skips disabled ones)
  SERVICES=$(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
     true)

  if [ -z "$SERVICES" ]; then
    # Fallback to common names
    SERVICES="Wi-Fi Ethernet Thunderbolt Ethernet USB 10/100 LAN"
  fi

  RESET_COUNT=0
  while IFS= read -r service; do
    # Strip leading asterisk (disabled services)
    service_clean=$(echo "$service" | sed 's/^\*//')
    [ -z "$service_clean" ] && continue

    # Check if this service has a hardware port (skip VPN/service-only entries)
    if networksetup -listallhardwareports 2>> output.log | grep -A1 "Port: $service_clean$" | grep -q "Device:"
      || [ "$service_clean" = "Wi-Fi" ]; then
      (
        networksetup -setsearchdomains "$service_clean" "Empty" 2>> output.log
        # Set to "empty" so macOS uses DHCP-assigned DNS (not hardcoded 1.1.1.1)
        sudo networksetup -setdnsservers "$service_clean" "empty" 2>> output.log
      ) && RESET_COUNT=$((RESET_COUNT + 1)) || true
    fi
  done <<< "$SERVICES"

  ok "DNS reset on $RESET_COUNT service(s)"
fi

#  3. Disable rogue proxy settings (default only  gateway uses proxy in non-exit-node path)
if [ "$POST_WAKE" = "false" ]; then
  step "Disabling proxy settings (SOCKS, web, secure web)"
  for svc in $(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
     echo "Wi-Fi"); do
    svc_clean=$(echo "$svc" | sed 's/^\*//')
    [ -z "$svc_clean" ] && continue
    disable_all_proxies "$svc_clean"
  done
  ok "Proxy settings cleared"
else
  step "Proxy settings: leaving untouched (post-wake keeps gateway SOCKS wiring alive)"
```

```
fi

#  4. Flush macOS DNS cache (always  safe and useful in both modes)
step "Flushing DNS cache"
if command -v dscacheutil >> output.log 2>&1; then
  flush_dns_cache
  ok "DNS cache flushed (mDNSResponder HUP)"
else
  warn "dscacheutil not found"
fi

#  5. Clear stale routing table (always  safe in both modes)
step "Flushing stale routing table entries"
# Resolve physical default gateway IP before flushing
# We cannot trust `route get default` here because Tailscale may have already hijacked
# the default route to utunX. We must read the actual DHCP router assignment from the hardware.
# PHYSICAL_IFACE and PHYSICAL_GW resolved upfront

# Instead of full route -n flush (which destroys macOS loopback/multicast and wedges the OS until reboot),
# we cleanly delete the Tailscale overrides and Cloudflare WARP bypass routes in ALL modes.
# This ensures tailscaled isn't blocked from reaching its control plane when waking from sleep.
sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
remove_warp_bypass_routes
ok "Routing overrides cleared"

if [ "$POST_WAKE" = "false" ]; then
  disable_killswitch
  ok "Routing table flush completed via targeted deletes and firewall disabled"
else
  ok "Routing overrides cleared (post-wake mode to preserve Colima bridge100)"
fi

# In post-wake mode we don't bounce Wi-Fi, so we MUST manually restore the physical default route
if [ "$POST_WAKE" = "true" ] && [ -n "$PHYSICAL_GW" ]; then
  # Ensure physical gateway is set just in case it was dropped
  sudo -n route add default "$PHYSICAL_GW" >> output.log 2>&1 || true
fi

# CONDITIONAL BYPASS ROUTE RESTORATION:
# If this was a post-wake recovery and the exit node is active, the route flush
# just wiped the static bypass routes for the WARP endpoints. We must re-add
# them immediately to prevent a routing loop when Tailscale starts sending traffic.
if [ "$POST_WAKE" = "true" ]; then
  if [ -z "${ACTIVE_IF:-}" ]; then
    ACTIVE_IF=$(route get default 2>> output.log | awk '/interface:/{print $2; exit}')
    if [ -z "$ACTIVE_IF" ] || [[ "$ACTIVE_IF" =~ ^utun ]] || [[ "$ACTIVE_IF" =~ ^tun ]]; then
      for i in en0 en1 en2 en3; do
        if ifconfig "$i" 2>> output.log | grep -q 'status: active'; then
          ACTIVE_IF="$i"; break
        fi
      done
    fi
    [ -z "$ACTIVE_IF" ] && ACTIVE_IF="en0"
  fi

  echo "Re-adding host bypass routes for Cloudflare WARP endpoints..."
  remove_warp_bypass_routes
  if [ -n "${PHYSICAL_GW:-}" ]; then
    add_warp_bypass_routes "$PHYSICAL_GW"
  else
    add_warp_bypass_routes "$ACTIVE_IF"
  fi

  # Also re-apply the default route pointing to the Tailscale interface
  ts_iface=$(ifconfig 2>> output.log | grep -B4 "inet 100." | grep -E '^[a-z0-9]+' | cut -d: -f1 | head -n 1)
  if [ -n "$ts_iface" ]; then
    echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
    # DO NOT delete the physical default route! It crashes Colima.
    # Use the /1 trick to mathematically override the default route.
    sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
    sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
    enable_killswitch
  fi
fi

#  6. Stop gateway Docker containers (default) / Force-recreate warp if unhealthy (post-wake)
if [ "$POST_WAKE" = "true" ]; then
  #  6a. Check for Roaming Network Mismatch (Bridged vs Enterprise Wi-Fi)
```

```
  if [ -f "/tmp/nullexit_colima_mode.txt" ] && [ -f "$SCRIPT_DIR/../.lan_p2p_detected" ]; then
    current_colima_mode=$(cat /tmp/nullexit_colima_mode.txt)
    p2p_allowed=$(cat "$SCRIPT_DIR/../.lan_p2p_detected")
    required_colima_mode="bridged"
    [ "$p2p_allowed" == "false" ] && required_colima_mode="shared"

    if [ "$current_colima_mode" != "$required_colima_mode" ]; then
      warn "Network roaming mismatch detected!"
      warn "Colima is running in '$current_colima_mode' mode but new network requires '$required_colima_mode'."
      warn "Executing full gateway restart to protect against MAC-spoofing deauthentication..."
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Roam] Mismatch ($current_colima_mode vs $required_colima_mode). Triggering toggle.sh --restart." >> output.log
      nohup bash "$SCRIPT_DIR/../toggle.sh" --restart >/dev/null 2>&1 &
      exit 0
    fi
  fi

  step "Checking Colima VM state"
  colima_status=$(colima status 2>&1)
  if echo "$colima_status" | grep -qi "running"; then
    ok "Colima VM is running"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima status check passed." >> output.log
  else
    warn "Colima VM is NOT running or unresponsive"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima is unhealthy or dead! Output: $colima_status" >> output.log
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Attempting to restart Colima..." >> output.log
    colima restart >> output.log 2>&1 || warn "Failed to restart Colima"
  fi

  step "Checking warp container health (gluetun UDP tunnel)"
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Inspecting warp container state..." >> output.log
  # The warp container is the most failure-prone link in the chain at wake/roam:
  # its UDP NAT binding to Cloudflare's edge (162.159.192.1:2408) silently dies
  # during a Wi-Fi roam. We verify the tunnel is actually passing traffic via curl.
  WARP_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container status: $WARP_STATE" >> output.log

  if [ "$WARP_STATE" = "missing" ]; then
    warn "warp container is missing  leaving gateway containers alone (full relaunch requires \"$SCRIPT_DIR/toggle.sh\")"
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container missing from Docker engine. Requires full toggle.sh launch." >> output.log
  else
    tmp_err=$(mktemp)
    if docker compose exec -T warp wget -qO- --timeout=3 https://www.cloudflare.com/cdn-cgi/trace 2>"$tmp_err" | grep -q "warp=on"; then
      ok "warp container is healthy and actively tunneling traffic (no recreate needed)"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container passed live traffic healthcheck (Cloudflare trace successful)." >> output.log
    else
      warn "warp container tunnel is dead (Docker status: $WARP_STATE)  force-recreating all gateway containers to restore network namespace"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container is DEAD or STUCK. Forcing recreation of gateway stack..." >> output.log
      if [ -s "$tmp_err" ]; then
        err_msg=$(cat "$tmp_err" | tr -d '\r')
        echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Healthcheck failure details: $err_msg" >> output.log
      fi
      docker compose up -d --force-recreate >> output.log 2>&1 || warn "force-recreate failed"

      NEW_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
      ok "warp container health after recreate: $NEW_STATE"
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container health after recreation: $NEW_STATE" >> output.log
    fi
    rm -f "$tmp_err"
  fi
else
  step "Stopping gateway Docker containers"
  if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
    if [ -f "docker-compose.yml" ]; then
      docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
    else
      warn "docker-compose.yml not found in current directory"
    fi
  else
    warn "Docker not available  skipping"
  fi
fi
```

```
611
612 #  6b. Stopping sleep prevention (caffeinate)  default only
613 if [ "$POST_WAKE" = "false" ]; then
614   step "Stopping sleep prevention"
615   stop_pidfile_daemon "$PID_CAFFEINATE" "system sleep prevention"
616 else
617   step "Sleep prevention: leaving untouched (caffeinate already re-acquired by parent shell)"
618 fi
619
620 #  6c. Stopping DNS Watcher  default only
621 if [ "$POST_WAKE" = "false" ]; then
622   step "Stopping DNS Watcher"
623   stop_pidfile_daemon "$PID_DNS_WATCHER" "background DNS Watcher"
624 else
625   step "DNS Watcher: leaving untouched (it keeps re-hijacking DNS on every roam)"
626 fi
627
628 #  7. Power-cycle Wi-Fi  default only
629 if [ "$POST_WAKE" = "false" ]; then
630   step "Power-cycling Wi-Fi"
631   bounce_wifi_interfaces
632 else
633   step "Wi-Fi: leaving untouched (post-wake keeps the current Wi-Fi association)"
634 fi
635
636 #  8. Verify (default checks direct internet; post-wake verifies the gateway)
637 if [ "$POST_WAKE" = "true" ]; then
638   step "Verifying gateway is healthy after post-wake refresh"
639
640   # Wait a moment for tailscale + warp healthcheck to settle
641   sleep 3
642
643   GATEWAY_OK=false
644   if command -v dig >> output.log 2>&1; then
645     # Source procedural ports if available
646     if [ -f "/tmp/nullexit-ports.env" ]; then
647       source "/tmp/nullexit-ports.env"
648     fi
649     DNS_PORT=${DNS_PROXY_PORT:-5354}
650
651     # AdGuard is exposed at localhost:${DNS_PORT}. If this resolves, the warp tunnel
652     # is functioning properly because AdGuard routes upstream to warp.
653     if dig +tcp @127.0.0.1 -p "$DNS_PORT" google.com +short +timeout=5 >> output.log 2>&1; then
654       ok "Gateway DNS (AdGuard via localhost:${DNS_PORT}) works"
655       GATEWAY_OK=true
656     else
657       warn "Gateway DNS not responding yet"
658     fi
659   fi
660
661   # Direct host  gateway check via Tailscale ping (relayed pong counts as success)
662   if command -v tailscale >> output.log 2>&1; then
663     TS_IP=""
664     [ -f .gateway_ip ] && TS_IP=$(read_adguard_ip || true)
665     if [ -n "$TS_IP" ] && tailscale ping --until-direct=false -c 1 --timeout 4s "$TS_IP" 2>> output.log | grep
      -q pong; then
666       ok "Gateway $TS_IP reachable via Tailscale"
667       GATEWAY_OK=true
668     else
669       warn "Tailscale ping to gateway did not pong  exit-node may still be re-establishing"
670     fi
671   fi
672
673   if [ "$GATEWAY_OK" = "true" ]; then
674     echo -e "  ${GREEN}${BOLD} Gateway is healthy after post-wake refresh.${NC}"
675   else
676     echo -e "  ${YELLOW}${BOLD} Gateway recovery is in-progress.${NC}"
677     echo "    The route may need another 30s before queries succeed."
678     echo "    If it stays broken for >60s, run: bash \"$SCRIPT_DIR/toggle.sh\""
679   fi
680 else
681   verify_internet_connectivity
682 fi
683
684 #  Summary
685 echo ""
686 echo -e "${BOLD}${NC}"
687 if [ "$POST_WAKE" = "true" ]; then
688   echo -e "${BOLD}POST-WAKE SUMMARY${NC}"
689 else
690   echo -e "${BOLD}RECOVERY SUMMARY${NC}"
```

```
fi
echo -e "${BOLD}${NC}"

# Show final DNS state
FINAL_DNS=$(networksetup -getdnsservers "Wi-Fi" 2>/dev/null || echo "N/A")
echo "  DNS (Wi-Fi):  $FINAL_DNS"

FINAL_DNS2=$(networksetup -getdnsservers "Ethernet" 2>/dev/null || true)
if [[ -z "$FINAL_DNS2" || "$FINAL_DNS2" == *"not a recognized"* || "$FINAL_DNS2" == *"Error"* ]]; then
  FINAL_DNS2="N/A (Not configured)"
fi
echo "  DNS (Ethernet): $FINAL_DNS2"

echo ""

if [ "$POST_WAKE" = "false" ]; then
  if [ "$INTERNET_OK" = "true" ]; then
    echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
  else
    echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
    echo "    This might mean:"
    echo "    - Your Wi-Fi is disconnected from the router"
    echo "    - Tailscale is still interfering (try restarting tailscaled:"
    echo "        brew services restart tailscale)"
    echo "    - You need to re-connect to Wi-Fi in the menu bar"
    echo "    - Try toggling Wi-Fi off/on in the menu bar"
    echo ""
    echo "    Or try manually:"
    echo "      networksetup -setdnsservers Wi-Fi empty"
    echo "      tailscale down"
    echo "      sudo route delete -net 0.0.0.0/1"
    echo "      sudo route delete -net 128.0.0.0/1"
    echo "      brew services restart tailscale"
  fi
fi

# Reset sharing services to prevent AirDrop freezes after interface changes
# This is the SAME step in BOTH modes (post-wake inherits the previous fix from 10.27).
echo -e "  Resetting macOS sharing services (AirDrop/AirPlay)..."
reset_sharing_services

echo ""
if [ "$POST_WAKE" = "true" ]; then
  echo -e "${YELLOW}${BOLD}Post-wake refresh complete.${NC}"
  # Remove the WARP Watcher inhibit marker now that the gateway is fully healthy.
  # The watcher will resume normal liveness monitoring on its next 5s poll.
  rm -f "$WARP_INHIBIT_FILE"
  echo "[$(date -u +%FT%TZ)] POST-WAKE complete  WARP Watcher inhibit released." >> output.log
else
  echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
fi
echo ""
```

8 setup.sh

```
#!/bin/bash
# setup.sh  nullexit one-shot setup script (macOS)
# Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
set -euo pipefail

# Source common formatting and helper functions
source "$(dirname "$0")/scripts/common.sh"

# Run common installation logic
source "$(dirname "$0")/scripts/setup-common.sh"

#  Compile macOS Shortcuts
echo -e "\n${BOLD}Compiling macOS Desktop Shortcuts...${NC}"
if command -v osacompile &> /dev/null; then
    osacompile -o "Toggle Gateway.app" "Toggle-Gateway.applescript" >/dev/null 2>&1
    osacompile -o "Recover Gateway.app" "Recover-Gateway.applescript" >/dev/null 2>&1
    echo "   Created 'Toggle Gateway.app' and 'Recover Gateway.app'"
fi

#  Done
echo ""
```

```
echo -e "${GREEN}${BOLD}${NC}"
echo -e "${GREEN}${BOLD}            Setup complete!                  ${NC}"
echo -e "${GREEN}${BOLD}${NC}"
echo ""
echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
echo "   https://login.tailscale.com/admin/machines"
echo "   Find your gateway  '...'  'Edit route settings'  enable Exit Node"
echo ""

echo -e "${YELLOW}${BOLD}Sudoers / Background Execution Requirements:${NC}"
echo "  To allow silent, non-interactive execution of the firewall and network routes (especially during sleep/wake recovery), you must configure passwordless sudo."
echo "  Run this exact command in your terminal (grants are scoped to least privilege):"
echo "    echo \"\$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /sbin/route, /sbin/ifconfig, /usr/bin/dscacheutil -flushcache, /usr/bin/killall -HUP mDNSResponder, /usr/bin/killall sharingd rapportd, /usr/bin/pkill -9 tailscaled, /usr/bin/pkill -f dns-proxy.py, /bin/kill, /usr/bin/true, /opt/homebrew/bin/brew services * tailscale, /opt/homebrew/bin/brew uninstall tailscale, /usr/local/bin/brew services * tailscale, /usr/local/bin/brew uninstall tailscale, /usr/bin/python3 -I \$PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I \$PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I \$PWD/scripts/dns-proxy.py\" | sudo tee /etc/sudoers.d/nullexit >/dev/null && sudo visudo -cf /etc/sudoers.d/nullexit"
echo ""

if [[ -n "${TS_IP:-}" ]]; then
    echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
    echo ""
fi

echo -e "${BOLD}To toggle the gateway on/off:${NC}"
echo "  macOS:   double-click the 'Toggle Gateway' app icon!"
echo "  Windows: double-click Toggle-Gateway.bat"
echo ""

if [[ "$OSTYPE" == "darwin"* ]]; then
    # LuLu localhost filtering check
    if [ -d "/Applications/LuLu.app" ] || systemextensionsctl list 2>/dev/null | grep -q "com.objective-see.lulu"; then
        echo -e "${YELLOW}${BOLD}LuLu Firewall Detected:${NC}"
        echo "  LuLu's localhost (loopback) filtering is OFF by default. To properly protect the"
        echo "  local nullexit proxy ports from malicious local processes, you must manually enable"
        echo "  'Allow Loopback' (or block it selectively) in LuLu's Preferences -> Rules."
        echo ""
    fi

    echo -e "${YELLOW}${BOLD}Note on macOS Tailscale Sandboxing:${NC}"
    echo "  The Tailscale GUI app (tailscale-app) installed on macOS is sandboxed."
    echo "  While you cannot run a Tailscale SSH server on this Mac host, you can"
    echo "  still SSH outbound to other devices on the mesh using their MagicDNS"
    echo "  addresses directly (e.g. ssh user@device.tailnet-name.ts.net)."
    echo ""
fi
```

9 scripts/setup-linux.sh

```
#!/bin/bash
# scripts/setup-linux.sh  nullexit one-shot setup script (Linux)
# Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
set -euo pipefail

# Source common formatting and helper functions
source "$(dirname "$0")/common.sh"

# Run common installation logic
source "$(dirname "$0")/setup-common.sh"

#  Compile Linux Desktop Shortcuts
echo -e "\n${BOLD}Creating Linux Desktop Shortcuts...${NC}"
DIR="$(pwd)"
cat <<EOF > "Toggle Gateway.desktop"
[Desktop Entry]
Version=1.0
Name=Toggle Gateway
Comment=Start or Stop Nullexit Gateway
Exec=bash -c "cd '$DIR' && ./toggle.sh; read -p 'Press Enter to close...' dummy"
Icon=utilities-terminal
Terminal=true
Type=Application
```

```
Categories=Network;Security;
EOF

cat <<EOF > "Recover Gateway.desktop"
[Desktop Entry]
Version=1.0
Name=Recover Gateway
Comment=Emergency DNS and Network Recovery
Exec=bash -c "cd '$DIR' && ./recover.sh; read -p 'Press Enter to close...' dummy"
Icon=utilities-terminal
Terminal=true
Type=Application
Categories=Network;Security;
EOF
echo "   Created 'Toggle Gateway.desktop' and 'Recover Gateway.desktop'"

#  Done
echo ""
echo -e "${GREEN}${BOLD}${NC}"
echo -e "${GREEN}${BOLD}            Setup complete!                  ${NC}"
echo -e "${GREEN}${BOLD}${NC}"
echo ""
echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
echo "   https://login.tailscale.com/admin/machines"
echo "   Find your gateway  '...'  'Edit route settings'  enable Exit Node"
echo ""

if [[ -n "${TS_IP:-}" ]]; then
    echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
    echo ""
fi

echo -e "${BOLD}To toggle the gateway on/off:${NC}"
echo "  Linux: double-click the 'Toggle Gateway' desktop icon!"
echo "  (or run ./toggle.sh in terminal)"
echo ""
```

10 docker-compose.yml

```
services:
  warp:
    image: qmcgaw/gluetun:v3.41.1
    container_name: warp
    hostname: nullexit-gateway
    cap_add:
      - NET_ADMIN
    devices:
      - /dev/net/tun:/dev/net/tun
    ports:
      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/udp"    # AdGuard DNS: host binds 5354 (the mapped DNS-proxy port); 5335 is AdGuard's in-container port
      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/tcp"    # AdGuard DNS (TCP)
      - "41642:41642/udp"  # Tailscale WireGuard (direct connection to host)
      - "127.0.0.1:${SOCKS_PROXY_PORT:-1080}:1080/tcp"    # SOCKS5 proxy (routed through WARP tunnel)
      - "127.0.0.1:${TOR_SOCKS_PORT:-9050}:9050/tcp"  # Tor SOCKS5 proxy (procedurally generated)
      - "127.0.0.1:${TOR_CONTROL_PORT:-9051}:9051/tcp"  # Tor Control port (procedurally generated)
    environment:
      - TZ=America/New_York
      - VPN_SERVICE_PROVIDER=custom
      - VPN_TYPE=wireguard
      - WIREGUARD_PRIVATE_KEY=${WIREGUARD_PRIVATE_KEY}
      - WIREGUARD_PUBLIC_KEY=${WIREGUARD_PUBLIC_KEY}
      - WIREGUARD_ADDRESSES=${WIREGUARD_ADDRESSES}
      - WIREGUARD_ENDPOINT_IP=${WARP_ENDPOINT_1:-162.159.192.1}
      - WIREGUARD_ENDPOINT_PORT=2408
      # Set to 25s (WireGuard default) to beat standard 30s hardware NAT/firewall UDP timeouts
      - WIREGUARD_PERSISTENT_KEEPALIVE_INTERVAL=25s
      # MTU Fix (Critical): Prevent packet fragmentation from double-encryption
      - WIREGUARD_MTU=1280
      # Give WARP more time to establish the connection before internal restart (default is 6s which is too short)
      - HEALTH_VPN_DURATION_INITIAL=30s
      # Use plaintext DNS resolvers instead of DNS-over-TLS (DoT) to prevent startup failures on networks where port 853 is blocked
      - DNS_UPSTREAM_RESOLVER_TYPE=plain
      - DNS_UPSTREAM_PLAIN_ADDRESSES=1.1.1.1:53,8.8.8.8:53
```

```
      # Allow local network traffic to bypass WARP so Tailscale can establish fast, direct peer-to-peer connections (no DERP relays) on the LAN
      - FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8
      # Built-in DNS blocking for ads, malicious domains, and tracking
      # NOTE: This is the massive "Hagezi-style" built-in blocklist.
      # You can turn these to 'on' for maximum security if you have spare RAM
      # or are running this on a cloud exit node with adequate autonomous resources (requires approx 1-2GB of RAM).
      # Disabled intentionally: AdGuard Home provides the DNS filtering layer instead.
      # If AdGuard fails, there is no fallback filtering, but the healthcheck
      # dependency chain ensures Tailscale won't come up if AdGuard is down.
      - BLOCK_MALICIOUS=off
      - BLOCK_SURVEILLANCE=off
      - BLOCK_ADS=off
    dns:
      - 1.1.1.1
    sysctls:
      - net.ipv4.ip_forward=1
    volumes:
      - ./post-rules.txt:/iptables/post-rules.txt:ro
    restart: unless-stopped
    healthcheck:
      test: ["CMD-SHELL", "/gluetun-entrypoint healthcheck"]
      interval: 2s
      timeout: 5s
      retries: 15
      start_period: 60s
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

  tailscale:
    image: tailscale/tailscale:v1.98.4
    container_name: tailscale
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
      adguardhome:
        condition: service_healthy
    cap_add:
      - NET_ADMIN
      - NET_RAW
      - SYS_MODULE
    sysctls:
      - net.ipv4.ip_forward=1
      - net.ipv6.conf.all.forwarding=1
    environment:
      - TZ=America/New_York
      - TS_EXTRA_ARGS=--advertise-exit-node
      - PORT=41642
      - TS_USERSPACE=false
      - TS_AUTHKEY=${TS_AUTHKEY}
      - TS_STATE_DIR=/var/lib/tailscale
      # (Note: For headless/cloud deployments, you can bypass this footgun by
      # using Tailscale Ephemeral Keys which auto-clean themselves).
      - TS_AUTH_ONCE=true
      # MTU Optimization: Force Tailscale packets to be smaller than WARP's 1280 MTU
      # to prevent WireGuard-in-WireGuard fragmentation, completely eliminating bufferbloat
      - TS_DEBUG_MTU=1200
      # Enable Tailscale's built-in SOCKS5 proxy to route out of the WARP tunnel
      - TS_SOCKS5_SERVER=:1080
    devices:
      - /dev/net/tun:/dev/net/tun
    volumes:
      - tailscale_state:/var/lib/tailscale
    restart: unless-stopped
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

  routing-fix:
    image: nullexit-routing-fix:v1.0.0
    build:
      context: ./scripts
      dockerfile: Dockerfile.routing-fix
    container_name: routing-fix
```

```
    network_mode: service:warp
    # NOTE: Deliberately NOT sharing tailscale's PID namespace  the shared
    # namespace caused routing-fix to be killed whenever tailscale restarted
    # (e.g. during gluetun health-check flaps), which dropped all routing
    # table changes and left traffic leaking through eth0 instead of tun0.
    cap_add:
      - NET_ADMIN
      - SYSLOG
    depends_on:
      warp:
        condition: service_healthy
      # rule-compiler dependency removed: it is now a hash-gated one-off run by
      # toggle.sh before `up` (15.8.8). This service boots on the existing
      # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
    environment:
      - TZ=America/New_York
      - WARP_ENDPOINT=${WARP_ENDPOINT_1:-162.159.192.1}
      - BLOCKED_COUNTRIES=${BLOCKED_COUNTRIES}
      - TAILSCALE_ALLOW_LAN_P2P=${TAILSCALE_ALLOW_LAN_P2P:-false}
    volumes:
      - ./scripts/routing-fix.sh:/routing-fix.sh:ro
      # SECURITY: ./adguard/work/ is bind-mounted into this container.
      # Never write to it from the host or another container. The
      # single allowed writer is `rule-compiler` container  manual
      # edits corrupt the AdGuard cache + BoltDB session store.
      - ./adguard/work/userfilters:/userfilters:ro
      - ./adguard/work:/adguard_work:ro
      - .:/app
    restart: unless-stopped
    stop_grace_period: 1s
    logging:
      driver: "json-file"
      options:
        max-size: "1m"
        max-file: "1"

  adguardhome:
    # WARNING: AdGuard is pre-configured with hardcoded credentials (admin/nullexit).
    # This is fine for a local Mac running behind a NAT, but if you are deploying
    # this on a public cloud VPS with port 80/3000 exposed, change this immediately!
    image: adguard/adguardhome:v0.107.77
    container_name: adguardhome
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
      # rule-compiler dependency removed: it is now a hash-gated one-off run by
      # toggle.sh before `up` (15.8.8). This service boots on the existing
      # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
    environment:
      - TZ=America/New_York
    healthcheck:
      test: ["CMD-SHELL", "nc -z 127.0.0.1 5335 || exit 1"]
      interval: 2s
      timeout: 5s
      retries: 10
    volumes:
      # SECURITY: ./adguard/work/ is bind-mounted into this container.
      # Never write to it from the host or another container. The
      # single allowed writer is `rule-compiler` container  manual
      # edits corrupt the AdGuard cache + BoltDB session store.
      - ./adguard/work:/opt/adguardhome/work
      - ./adguard/conf:/opt/adguardhome/conf
    restart: unless-stopped
    stop_grace_period: 2s
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

  rule-compiler:
    image: nullexit-rule-compiler:v1.0.0
    build:
      context: ./scripts
      dockerfile: Dockerfile.rule-compiler
    container_name: rule-compiler
    # `compile` profile: excluded from `docker compose up`. toggle.sh runs it
    # on-demand via `docker compose run --build --rm rule-compiler`, hash-gated on
    # black_list.txt/white_list.txt, so ~340k rules are re-deduped only when a list
    # actually changes instead of every toggle (~28s saved on the common path, 15.8.8).
```

```
    profiles: ["compile"]
    volumes:
      - .:/app

  tor:
    build: ./docker/tor
    image: nullexit-tor:v1.0.0
    container_name: tor
    network_mode: service:warp
    depends_on:
      warp:
        condition: service_healthy
    environment:
      - TZ=America/New_York
      - TOR_USE_BRIDGES=${TOR_USE_BRIDGES:-false}
      - TOR_BRIDGE_LINES_FILE=${TOR_BRIDGE_LINES_FILE:-./tor-bridges.txt}
      - TOR_PASSWORD=${ADGUARD_PASSWORD:-nullexit}
      - TOR_EXCLUDE_EXIT_NODES=${TOR_EXCLUDE_EXIT_NODES:-{us}}
    working_dir: /nullexit
    volumes:
      - ./:/nullexit:ro
      - ./output.log:/output.log:rw
    restart: unless-stopped
    logging:
      driver: "json-file"
      options:
        max-size: "10m"
        max-file: "3"

volumes:
  tailscale_state:

networks:
  default:
    ipam:
      config:
        - subnet: 10.200.1.0/24
```

11 scripts/common.sh

```
#!/bin/bash
# common.sh  shared bash utilities for nullexit scripts

#  Formatting
RED='\033[0;31m'
GREEN='\033[0;32m'
YELLOW='\033[1;33m'
BLUE='\033[0;34m'
BOLD='\033[1m'
NC='\033[0m'

#  Constants & Environment
export MARKER_FILE="/tmp/nullexit-gateway-active.marker"
export PID_CAFFEINATE="/tmp/nullexit-caffeinate.pid"
export PID_DNS_WATCHER="/tmp/nullexit-dns-watcher.pid"
export DEFAULT_WARP_ENDPOINT_1="162.159.192.1"
export DEFAULT_WARP_ENDPOINT_2="162.159.193.1"

setup_standard_path() {
  export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
}

step() { echo -e "\n[$(date '+%Y-%m-%d %H:%M:%S')] ${BLUE}${BOLD} $*${NC}"; }
ok()   { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${GREEN} $*${NC}"; }
warn() { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${YELLOW} $*${NC}"; }
fail() { echo -e "[$(date '+%Y-%m-%d %H:%M:%S')]   ${RED} $*${NC}"; }
die()  { echo -e "\n[$(date '+%Y-%m-%d %H:%M:%S')]   ${RED} $*${NC}\n"; exit 1; }

#  Helper Functions

# Append a timestamped line to output.log (best-effort; never fails the caller).
# Used to record lifecycle breadcrumbs (EXEC/EXIT markers, phase changes) so a
# START/STOP that is killed or window-closed still leaves a retraceable trail.
log_line() {
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] $*" >> "${LOG_FILE:-output.log}" 2>/dev/null || true
}
```

```
37
38 # Run a lifecycle command with FULL observability: log the command, stream its
39 # combined stdout+stderr to the terminal AND append it to output.log via tee,
40 # then log the exit code explicitly. Returns the command's real exit code (not
41 # tee's). This closes the blind spot where `docker compose up` / `colima start`
42 # printed only to the terminal  so when they fail (or the run is interrupted),
43 # output.log shows exactly what the last command was and whether it succeeded.
44 #
45 # Usage:  run_logged docker compose up -d --build
46 run_logged() {
47   local logf="${LOG_FILE:-output.log}"
48   log_line "EXEC: $*"
49   # PIPESTATUS[0] is the command's exit code; tee (PIPESTATUS[1]) is ignored.
50   "$@" 2>&1 | tee -a "$logf"
51   local rc=${PIPESTATUS[0]}
52   log_line "EXIT $rc: $*"
53   return "$rc"
54 }
55
56 # Pure-bash timeout (no dependency on GNU coreutils' timeout)
57 # Runs a command with a safety cutoff so a wedged daemon can't hang the script.
58 run_with_timeout() {
59   local timeout_sec="$1"
60   shift
61   if [ $# -eq 0 ]; then return 1; fi
62
63   "$@" &
64   local cmd_pid=$!
65   CURRENT_BG_PID="$cmd_pid"
66
67   (
68     sleep "$timeout_sec"
69     if kill -0 "$cmd_pid" 2>> output.log; then
70       echo -e "\n[Timeout] Command '$*' exceeded $timeout_sec seconds. Terminating..." >&2
71       kill -15 "$cmd_pid" 2>> output.log || true
72       sleep 2
73       if kill -0 "$cmd_pid" 2>> output.log; then
74         kill -9 "$cmd_pid" 2>> output.log || true
75       fi
76     fi
77   ) &
78   local watcher_pid=$!
79
80   # Disable set -e temporarily to capture exit status of wait
81   set +e
82   wait "$cmd_pid" 2>> output.log
83   local exit_status=$?
84   set -e
85
86   # Kill the watcher since the command finished
87   kill "$watcher_pid" 2>> output.log && wait "$watcher_pid" 2>> output.log || true
88
89   CURRENT_BG_PID=""
90   return "$exit_status"
91 }
92
93 # Start Colima and return the moment Docker is actually reachable, instead of
94 # blocking on colima 0.10.3's post-provision hang. Observed on macOS/vz: the VM
95 # reaches READY and creates the `colima` docker context in ~10-15s, then
96 # `colima start` frequently fails to return for ~110s (until a watchdog kills it)
97 # even though Docker is fully usable the entire time. We poll the colima docker
98 # context and proceed as soon as it answers, then reap the (usually hung) starter.
99 # The VM keeps running and `colima status` stays accurate. $@ = colima start args.
100 # Returns 0 once Docker responds, 1 if it never does within the cap.
101 colima_start_until_ready() {
102   local max_wait=90 waited=0 ready=false start_pid
103   echo "[$(date '+%Y-%m-%d %H:%M:%S')] EXEC(bg): colima start $*" >> output.log
104   colima start "$@" >> output.log 2>&1 &
105   start_pid=$!
106   while [ "$waited" -lt "$max_wait" ]; do
107     if run_with_timeout 5 docker --context colima info >/dev/null 2>&1; then
108       ready=true
109       break
110     fi
111     # If colima start exited on its own (genuine finish or hard failure), stop waiting.
112     if ! kill -0 "$start_pid" 2>/dev/null; then
113       run_with_timeout 5 docker --context colima info >/dev/null 2>&1 && ready=true
114       break
115     fi
116     sleep 3
117     waited=$((waited + 3))
```

```
  done
  # Reap the starter (harmless if it already exited); this does NOT stop the VM.
  kill "$start_pid" 2>/dev/null || true
  wait "$start_pid" 2>/dev/null || true
  if [ "$ready" = "true" ]; then
    docker context use colima >/dev/null 2>&1 || true
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker ready after ~${waited}s (starter reaped)." >> output.log
    return 0
  fi
  echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker NOT ready after ${max_wait}s." >> output.log
  return 1
}

# Hash the files that feed the long-running locally-built images that `docker
# compose up` starts (routing-fix, tor). toggle.sh uses this to skip
# `docker compose --build` when nothing changed: BuildKit re-evaluating those
# images on the 600MB Colima VM costs ~30s per toggle even when every layer is
# CACHED (15.8.8). Includes the Dockerfiles, everything they COPY (routing-fix.sh,
# tor/entrypoint.sh, the logger/ Go source), and docker-compose.yml. The
# rule-compiler image is NOT here  it is a `compile`-profiled one-off, built+run
# by the gated compile step (15.8.8). Prints a sha256 hex digest, or empty.
compute_build_inputs_hash() {
  local root="${SCRIPT_DIR:-.}" f
  {
    for f in \
      "$root/scripts/Dockerfile.routing-fix" \
      "$root/scripts/routing-fix.sh" \
      "$root/docker/tor/Dockerfile" \
      "$root/docker/tor/entrypoint.sh" \
      "$root/docker-compose.yml"; do
      printf '::%s::' "$f"; cat "$f" 2>/dev/null
    done
    find "$root/scripts/logger" -type f 2>/dev/null \
      | LC_ALL=C sort | while IFS= read -r f; do printf '::%s::' "$f"; cat "$f" 2>/dev/null; done
  } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
}

# Hash the ad-block lists the user controls: black_list.txt (custom blocks) and
# white_list.txt (the allow-list that governs what DNS resolves through). toggle.sh
# gates the ~28s rule recompile on this  editing either list changes the digest
# and triggers a one-off recompile of compiled_rules.txt + ip_blocklist.ipset;
# an unchanged toggle skips it entirely (15.8.8). Prints a sha256 hex digest.
compute_rules_hash() {
  local root="${SCRIPT_DIR:-.}"
  { printf '::black::'; cat "$root/black_list.txt" 2>/dev/null
    printf '::white::'; cat "$root/white_list.txt" 2>/dev/null
  } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
}

# Check if the gateway is active (either containers are running, or host DNS is hijacked)
is_gateway_active() {
  # 1. Check if containers are running (suppress stderr to avoid errors if docker is down)
  if run_with_timeout 15 docker compose ps --status running 2>/dev/null | grep -q 'warp'; then
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (warp container is running)" >> "$LOG_FILE"
    return 0
  fi

  # 2. Check if host DNS was hijacked (not 1.1.1.1 and not default/empty)
  local current_dns=""
  local active_svc=""
  if [[ "$OSTYPE" == "darwin"* ]]; then
    active_svc=$(get_active_service)
    current_dns=$(networksetup -getdnsservers "$active_svc" 2>> output.log || true)
  else
    current_dns=$(resolvectl dns 2>/dev/null | awk '/Global|Link/ {for(i=4;i<=NF;i++) print $i}' | head -n1)
  fi

  if [[ -n "$current_dns" && "$current_dns" != "1.1.1.1" && ! "$current_dns" =~ "There aren't any DNS Servers"
     ]]; then
    echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (DNS was hijacked: $current_dns)" >> "
     $LOG_FILE"
    return 0
  fi

  # 3. Check if SOCKS proxy is enabled (macOS only)
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local socks_proxy
    socks_proxy=$(networksetup -getsocksfirewallproxy "$active_svc" 2>> output.log || true)
    if echo "$socks_proxy" | grep -q "Enabled: Yes"; then
      echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (SOCKS proxy enabled)" >> "$LOG_FILE"
      return 0
```

```
    fi
  fi

  echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: FALSE (Containers down, DNS clean, SOCKS disabled)" >> "$LOG_FILE"
  return 1
}

# Decide whether the gateway SHOULD be running, from persisted *intent* rather
# than a live health probe. This is the correct signal for toggle.sh's
# STOP-vs-START decision: a disable should tear down whatever the last START
# left behind  it should NOT depend on whether Docker/Colima happens to be
# reachable right now.
#
# Why this exists: using is_gateway_active() for the decision means every toggle
# (including *disable*) runs `docker compose ps` first. When Colima is down the
# socket is absent, that call blocks for the full run_with_timeout window, and
# under `set -e` + the ERR trap a non-zero result at the `if  then  else  fi`
# boundary could abort the script BEFORE the START body runs  so the very path
# that would boot Colima never executed (observed as `EXIT: code=1 phase=start`).
#
# Signal source: MARKER_FILE is written at the end of a successful START
# (write_gateway_active_marker) and cleared at the top of STOP and in
# cleanup_handler. Reading it is instant and cannot hang or abort.
#
# ECHOES "running" or "stopped" to stdout (and ALWAYS returns 0). Callers read
# the string:  [ "$(gateway_state)" = "running" ] &&
#
# CRITICAL  why this echoes instead of using a 0/1 return code: on macOS
# /bin/bash 3.2, with `set -e` + an `ERR` trap + `set -x` (DEBUG_TRACE=true) all
# active, a function that does `return 1` triggers a SPURIOUS `set -e` abort in
# the caller even when the call is guarded (`if f; then`, or even `f || rc=$?`).
# The `return N` from the function, traced by `set -x`, trips the ERR trap. This
# is a genuine bash-3.2 interaction (see devref 15.11.8). Echoing a result and
# always returning 0 sidesteps `return`/`set -e`/`set -x` entirely  the function
# can never contribute a non-zero status, so no caller can be aborted by it.
gateway_state() {
  # Primary: persisted intent. Marker present => a START completed and no STOP
  # has cleared it yet.
  #
  # NOTE: toggle.sh runs a "defensive stale-marker clear" near the top of every
  # invocation (before this decision) to stop the wake-watcher firing against a
  # crashed gateway. That would wipe the marker before we read it, so toggle.sh
  # captures the marker's existence into GATEWAY_MARKER_AT_STARTUP *before* that
  # clear and we honor it here. If the variable is unset (other callers, or the
  # --restart pre-check which runs before the defensive clear), fall back to the
  # live marker file.
  if [ -n "${GATEWAY_MARKER_AT_STARTUP:-}" ]; then
    if [ "$GATEWAY_MARKER_AT_STARTUP" = "yes" ]; then echo "running"; return 0; fi
  elif [ -f "$MARKER_FILE" ]; then
    echo "running"; return 0
  fi

  # Marker absent. Normally this means "stopped". But cover the edge case where
  # the marker was lost (e.g. /tmp cleared, or gateway started by an older build)
  # yet containers are genuinely up  reconcile via is_gateway_active, but ONLY
  # when Docker is actually reachable, so a dead-socket probe can never hang. On
  # macOS Docker runs in the Colima VM; if that socket is absent, the gateway
  # definitively is not running.
  local docker_up="no"
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if [ -S /var/run/docker.sock ] || [ -S "$HOME/.colima/default/docker.sock" ]; then docker_up="yes"; fi
  else
    if [ -S /var/run/docker.sock ]; then docker_up="yes"; fi
  fi

  if [ "$docker_up" = "yes" ] && is_gateway_active; then
    echo "running"; return 0
  fi
  echo "stopped"
  return 0
}

# Reads an environment variable from a file (default: .env), strips quotes and whitespace
read_env_var() {
  local var_name="$1"
  local env_file="${2:-.env}"
  grep -E "^${var_name}=" "$env_file" 2>/dev/null | cut -d'=' -f2- | tr -d "\"\\'" | sed -e 's/^[[:space:]]*//' -e 's/[[:space:]]*$//' || echo ""
}

```

```
# Load KILL_SWITCH variable from .env (defaults to false if not found/empty)
export KILL_SWITCH
KILL_SWITCH=$(read_env_var "KILL_SWITCH" 2>/dev/null | tr '[:upper:]' '[:lower:]')
if [ -z "$KILL_SWITCH" ]; then
  KILL_SWITCH="false"
fi

# Function to get the active network service name (e.g., "Wi-Fi" or "USB 10/100 LAN")
get_active_service() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local iface
    iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')

    # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
    if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
      for i in en0 en1 en2 en3; do
        if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "
     inet "; then
          iface="$i"
          break
        fi
      done
    fi

    # Default fallback if still empty or not enX
    if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
      iface="en0"
    fi

    # Map interface (e.g., en0) to service name (e.g., Wi-Fi)
    local service
    service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B 1 "Device: $iface" | head -n 1 |
     sed -E 's/^\([0-9\*]+\) //')

    if [ -n "$service" ]; then
      echo "$service"
    else
      echo "Wi-Fi"
    fi
  else
    # Get the interface used for default internet routing
    local iface
    iface=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
    if [ -z "$iface" ]; then
      echo ""
      return 1
    fi
    echo "$iface"
  fi
}

get_en0_service() {
  local service
  service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B1 "Device: en0" | head -1 | sed -E 's
     /^\([0-9\*]+\) //' || true)
  if [ -z "$service" ]; then
    echo "Wi-Fi"
  else
    echo "$service"
  fi
}

get_warp_endpoint_1() { read_env_var "WARP_ENDPOINT_1" ".env" | grep -Eo '^[0-9\.]+' || echo "
     $DEFAULT_WARP_ENDPOINT_1"; }
get_warp_endpoint_2() { read_env_var "WARP_ENDPOINT_2" ".env" | grep -Eo '^[0-9\.]+' || echo "
     $DEFAULT_WARP_ENDPOINT_2"; }

restart_tailscaled_daemon() {
  echo "  [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
  echo "  [Tailscale Recovery] Attempting to restart tailscaled service..."

  if run_with_timeout 15 brew services restart tailscale >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled (user service)."
  elif run_with_timeout 15 sudo -n brew services restart tailscale >> output.log 2>&1; then
    echo "  [Tailscale Recovery] Successfully restarted tailscaled (system service)."
  else
    echo "  [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
  fi
  sleep 3
}

```

```
disconnect_tailscale_host() {
  local ts_bin="${1:-tailscale}"
  if command -v "$ts_bin" >> output.log 2>&1; then
    echo "  Disconnecting host Tailscale from mesh..."

    # Check if actually connected first (with timeout  tailscaled could be wedged)
    local status_ok=false
    if run_with_timeout 5 "$ts_bin" status >> output.log 2>&1; then
      status_ok=true
    else
      local status_exit=$?
      # If it was a quick exit code 1 (disconnected), it is still reachable
      if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
        status_ok=true
      fi
    fi

    if [ "$status_ok" = "false" ]; then
      restart_tailscaled_daemon
    fi

    # Explicitly reset any exit-node so it doesn't linger.
    "$ts_bin" up >> output.log 2>&1 || true
    if run_with_timeout 10 "$ts_bin" set --ssh=true --accept-dns=false --exit-node= >> output.log 2>&1; then
      echo "  [] Exit-node preference cleared."
    else
      echo "  [!] tailscale up didn't respond (tailscaled may be wedged)"
    fi

    local ts_args=""
    if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
    if run_with_timeout 10 "$ts_bin" down $ts_args >> output.log 2>&1; then
      echo "  [] Tailscale disconnected."
    else
      echo "  [!] tailscale down didn't respond"
    fi
  else
    echo "  [!] tailscale CLI not found  skipping"
  fi
}

stop_pidfile_daemon() {
  local pid_file="$1"
  local label="$2"
  if [ -f "$pid_file" ]; then
    local pid
    pid=$(cat "$pid_file")
    if kill -0 "$pid" 2>/dev/null; then
      echo "  Stopping $label (PID $pid)..."
      sudo -n kill "$pid" 2>> output.log || kill "$pid" 2>> output.log || true
      sleep 0.5
      if kill -0 "$pid" 2>/dev/null; then
        sudo -n kill -9 "$pid" 2>> output.log || kill -9 "$pid" 2>> output.log || true
      fi
    fi
    rm -f "$pid_file"
  fi
}

disable_all_proxies() {
  local svc="$1"
  if [ -n "$svc" ]; then
    sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
    sudo -n networksetup -setwebproxystate "$svc" off 2>> output.log || true
    sudo -n networksetup -setsecurewebproxystate "$svc" off 2>> output.log || true
  fi
}

flush_dns_cache() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    sudo -n dscacheutil -flushcache 2>> output.log || true
    sudo -n killall -HUP mDNSResponder 2>> output.log || true
  else
    resolvectl flush-caches 2>> output.log || true
  fi
}

wait_for_dhcp_settle() {
  # Resolve physical interface (Wi-Fi or Ethernet)
  local physical_iface
```

```
  physical_iface=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
    getline; print $2; exit}')
  [ -z "$physical_iface" ] && physical_iface="en0"

  echo -n "Waiting for physical network interface ($physical_iface) DHCP lease to settle..."
  local settled=false
  for attempt in {1..60}; do
    PHYSICAL_GW=$(ipconfig getpacket "$physical_iface" 2>> output.log | awk -F'[{}]' '/router /{print $2}')
    local local_ip
    local_ip=$(ifconfig "$physical_iface" 2>/dev/null | awk '/inet /{print $2}')

    if [ -n "$PHYSICAL_GW" ] && [[ "$PHYSICAL_GW" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
       [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
       [[ "$local_ip" != "169.254."* ]]; then
      echo " settled (Interface IP: $local_ip, Router IP: $PHYSICAL_GW)."
      settled=true
      break
    fi

    if [ "$attempt" -gt 15 ] && [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] &&
     [[ "$local_ip" != "169.254."* ]]; then
      local fallback_gw
      fallback_gw=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
      if [ -n "$fallback_gw" ] && [[ "$fallback_gw" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]]; then
        PHYSICAL_GW="$fallback_gw"
      fi
      echo " static/settled detected (Interface IP: $local_ip, Gateway IP: ${PHYSICAL_GW:-unknown})."
      settled=true
      break
    fi
    echo -n "."
    sleep 0.5
  done
  if [ "$settled" = "false" ]; then
    echo " timed out or no active link. Proceeding anyway."
  fi
}

reset_sharing_services() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    sudo -n killall sharingd rapportd 2>> output.log || true
  fi
}

read_adguard_ip() {
  if [ -f ".gateway_ip" ]; then
    tr -d '\r' < ".gateway_ip" | awk 'NR==1{print $1;exit}'
  fi
}

is_kill_switch_enabled() {
  [ "$KILL_SWITCH" = "true" ]
}

add_warp_bypass_routes() {
  local target="${1:-}"
  local ep1
  ep1=$(get_warp_endpoint_1)
  local ep2
  ep2=$(get_warp_endpoint_2)

  if [[ "$OSTYPE" == "darwin"* ]]; then
    local routing_arg=""
    local msg_via=""
    if [ -n "$target" ]; then
      if [[ "$target" =~ ^en[0-9]+$ || "$target" == "bridge"* ]]; then
        routing_arg="-interface $target"
        msg_via="interface $target"
      else
        routing_arg="$target"
        msg_via="gateway $target"
      fi
    else
      local gateway_ip
      gateway_ip=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
      if [ -z "$gateway_ip" ]; then
        local iface
        iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
        if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
          iface="en0"
        fi
```

```
        routing_arg="-interface $iface"
        msg_via="interface $iface"
      else
        routing_arg="$gateway_ip"
        msg_via="gateway $gateway_ip"
      fi
    fi

    echo -e "\nAdding host bypass routes for Cloudflare WARP endpoints via $msg_via..."
    sudo -n route delete -host "$ep1" 2>/dev/null || true
    sudo -n route delete -host "$ep2" 2>/dev/null || true
    sudo -n route add -host "$ep1" $routing_arg >> output.log 2>&1 || true
    sudo -n route add -host "$ep2" $routing_arg >> output.log 2>&1 || true

    echo -e "Adding host bypass routes for Tailscale control plane via $msg_via..."
    sudo -n route delete -net 192.200.0.0/24 2>/dev/null || true
    sudo -n route add -net 192.200.0.0/24 $routing_arg >> output.log 2>&1 || true

    echo -e "Adding host bypass routes for Tailscale DERP relays via $msg_via..."
    local derp_ips
    derp_ips=$(curl -s --connect-timeout 5 https://login.tailscale.com/derpmap/default | grep -oE '"IPv4"[[:space:]]*:[[:space:]]*"[0-9.]+"' | cut -d'"' -f4 | grep -E '^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$' || true)
    if [ -n "$derp_ips" ]; then
      echo "$derp_ips" > /tmp/nullexit-derp-ips.txt
      for ip in $derp_ips; do
        sudo -n route delete -host "$ip" 2>/dev/null || true
        sudo -n route add -host "$ip" $routing_arg >> output.log 2>&1 || true
      done
    fi
  else
    local gateway_ip="${target}"
    if [ -z "$gateway_ip" ]; then
      gateway_ip=$(ip route show default 2>/dev/null | awk '/default/ {print $3}')
    fi
    if [ -n "$gateway_ip" ]; then
      sudo ip route add "$ep1" via "$gateway_ip" >> output.log 2>&1 || true
      sudo ip route add "$ep2" via "$gateway_ip" >> output.log 2>&1 || true
    fi
  fi
}

remove_warp_bypass_routes() {
  local ep1
  ep1=$(get_warp_endpoint_1)
  local ep2
  ep2=$(get_warp_endpoint_2)

  if [[ "$OSTYPE" == "darwin"* ]]; then
    echo -e "\nRemoving host bypass routes for Cloudflare WARP endpoints..."
    sudo -n route delete -host "$ep1" >> output.log 2>&1 || true
    sudo -n route delete -host "$ep2" >> output.log 2>&1 || true
    echo -e "Removing host bypass routes for Tailscale control plane..."
    sudo -n route delete -net 192.200.0.0/24 >> output.log 2>&1 || true
    if [ -f /tmp/nullexit-derp-ips.txt ]; then
      echo -e "Removing host bypass routes for Tailscale DERP relays..."
      while read -r ip; do
        sudo -n route delete -host "$ip" >> output.log 2>&1 || true
      done < /tmp/nullexit-derp-ips.txt
      rm -f /tmp/nullexit-derp-ips.txt
    fi
  else
    sudo ip route del "$ep1" >> output.log 2>&1 || true
    sudo ip route del "$ep2" >> output.log 2>&1 || true
  fi
}

bounce_wifi_interfaces() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local wifi_port
    wifi_port=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print $2}')
    if [ -n "$wifi_port" ]; then
      sudo -n networksetup -setairportpower "$wifi_port" off 2>> output.log || true
      sleep 2
      sudo -n networksetup -setairportpower "$wifi_port" on 2>> output.log || true
      echo "Wi-Fi bounced (interface $wifi_port)"
      sleep 3
    else
      echo "Could not detect Wi-Fi interface. Bouncing fallback interfaces..."
      set +e
      for iface in en0 en1 en2 en3 en4 en5; do
```

```
        if ifconfig "$iface" >> output.log 2>&1; then
          sudo -n ifconfig "$iface" down 2>> output.log || true
          sudo -n ifconfig "$iface" up 2>> output.log || true
        fi
      done
      set -e
      echo "Fallback interfaces bounced"
    fi
  fi
}

get_active_interface() {
  local iface
  iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')

  # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
  if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
    for i in en0 en1 en2 en3; do
      if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "
     inet "; then
        iface="$i"
        break
      fi
    done
  fi

  # Default fallback if still empty or not enX
  if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
    iface="en0"
  fi

  echo "$iface"
}

enable_killswitch() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if ! is_kill_switch_enabled; then
      return 0
    fi

    local ext_if
    ext_if=$(get_active_interface)
    local ep1
    ep1=$(get_warp_endpoint_1)
    local ep2
    ep2=$(get_warp_endpoint_2)
    local pf_conf_path
    pf_conf_path="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/pf.conf"

    echo -e "\nEnabling macOS Packet Filter (PF) Kill-Switch on $ext_if..."

    # Enable PF and capture the reference token
    local token_out
    token_out=$(sudo -n pfctl -E 2>> output.log || true)
    local token
    token=$(echo "$token_out" | awk '/Token/{print $NF}')
    if [ -n "$token" ]; then
      echo "$token" > /tmp/nullexit-pf-token.txt
      echo "  [] PF enabled with token $token."
    else
      echo "  [!] WARNING: Could not capture PF reference token (may already be enabled or sudo failed)."
    fi

    # Load the rules into the anchor
    if sudo -n pfctl -D "ext_if=$ext_if" -a com.apple/nullexit -f "$pf_conf_path" >> output.log 2>&1; then
      echo "  [] Ruleset loaded into anchor com.apple/nullexit."
    else
      echo "  [!] ERROR: Failed to load ruleset into anchor."
    fi

    # Populate tables in the anchor
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T flush >> output.log 2>&1 || true
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep1" >> output.log 2>&1 || true
    sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$ep2" >> output.log 2>&1 || true

    if [ -f /tmp/nullexit-derp-ips.txt ]; then
      local derp_ips
      derp_ips=$(tr '\n' ' ' < /tmp/nullexit-derp-ips.txt)
      if [ -n "$derp_ips" ]; then
        sudo -n pfctl -a com.apple/nullexit -t derp_relays -T flush >> output.log 2>&1 || true
        sudo -n pfctl -a com.apple/nullexit -t derp_relays -T add $derp_ips >> output.log 2>&1 || true
```

```
      fi
    fi
    echo "  [] Kill-switch tables populated."
  fi
}

disable_killswitch() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if ! is_kill_switch_enabled; then
      return 0
    fi
    echo -e "\nDisabling macOS PF Kill-Switch..."

    # Flush rules from anchor com.apple/nullexit
    sudo -n pfctl -a com.apple/nullexit -F all >> output.log 2>&1 || true

    # Release the enable reference if token is present
    if [ -f /tmp/nullexit-pf-token.txt ]; then
      local token
      token=$(cat /tmp/nullexit-pf-token.txt)
      if [ -n "$token" ]; then
        sudo -n pfctl -X "$token" >> output.log 2>&1 || true
        echo "  [] Released PF enable reference token $token."
      fi
      rm -f /tmp/nullexit-pf-token.txt
    else
      echo "  [] Rules flushed from anchor com.apple/nullexit."
    fi
  fi
}


write_host_ips() {
  # Gather all active IPv4 addresses on the host and write to .host_ips
  # routing-fix.sh will read this file to explicitly whitelist them for direct Tailscale P2P
  local repo_dir
  repo_dir="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"

  if [[ "$OSTYPE" == "darwin"* ]]; then
    ifconfig | awk '/inet /{print $2}' | grep -v '127.0.0.1' > "$repo_dir/.host_ips" 2>/dev/null || true
  else
    ip -4 addr show 2>/dev/null | awk '/inet /{print $2}' | cut -d/ -f1 | grep -v '127.0.0.1' > "$repo_dir/.
     host_ips" 2>/dev/null || true
  fi
}

start_dns_watcher() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local target_ip=$1
    stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
    echo "  Starting background DNS Watcher for seamless Wi-Fi roaming..."
    nohup bash -c "
      source \"$SCRIPT_DIR/scripts/common.sh\"
      trap 'exit 0' SIGTERM SIGINT SIGHUP
      while true; do
        ACTIVE_IF=\$(get_active_service)
        if [ -n \"\$ACTIVE_IF\" ]; then
          CURRENT_DNS=\$(networksetup -getdnsservers \"\$ACTIVE_IF\" 2>/dev/null)
          if [ \"\$CURRENT_DNS\" != \"$target_ip\" ]; then
            networksetup -setdnsservers \"\$ACTIVE_IF\" \"$target_ip\" >/dev/null 2>&1
          fi
        fi
        sleep 30
      done
    " >> output.log 2>&1 &
    echo $! > "$DNS_WATCHER_PID_FILE"
  else
    local target_ip=$1
    stop_dns_watcher
    echo "  Starting background DNS Watcher for seamless Wi-Fi roaming..."
    nohup bash -c "
      trap 'exit 0' SIGTERM SIGINT SIGHUP
      while true; do
        if [ -n \"\$ACTIVE_SERVICE\" ]; then
          CURRENT_DNS=\$(resolvectl dns \"\$ACTIVE_SERVICE\" 2>/dev/null | grep -oE
     '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head -1)
          if [ \"\$CURRENT_DNS\" != \"$target_ip\" ]; then
            resolvectl dns \"\$ACTIVE_SERVICE\" \"$target_ip\" >/dev/null 2>&1 || true
          fi
        fi
        sleep 30
```

```
      done
    " >> output.log 2>&1 &
    echo $! > "$DNS_WATCHER_PID_FILE"
  fi
}

#  SOCKS5 Proxy (WARP Tunnel)
# When Tailscale's exit node can't route through WARP (due to userspace
# forwarding), we use a SOCKS5 proxy running inside the warp container's
# network namespace. Connections created by this proxy go through the
# kernel routing table (table 200 -> tun0 -> WARP), unlike Tailscale's
# userspace exit node which bypasses it.
#
# The SOCKS5 proxy is exposed on localhost:${SOCKS_PROXY_PORT} via Docker port mapping.
# `networksetup -setsocksfirewallproxy` on macOS routes system TCP traffic
# through the proxy, which then goes through WARP.
#
# IMPORTANT: CLI tools like curl do NOT respect macOS system proxy settings.
# They must use --socks5-hostname or the ALL_PROXY env variable explicitly.

enable_socks_proxy() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local svc="$1"
    if [ -z "$svc" ]; then
      svc="$ACTIVE_SERVICE"
    fi

    echo -n "  Enabling system-wide SOCKS5 proxy on $svc (localhost:$SOCKS_PROXY_PORT)... "
    sudo -n networksetup -setsocksfirewallproxy "$svc" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.log || true
    sudo -n networksetup -setsocksfirewallproxystate "$svc" on 2>> output.log || true

    # Also set on en0 service for macOS resolver consistency
    if [ "$svc" != "$EN0_SERVICE" ]; then
      sudo -n networksetup -setsocksfirewallproxy "$EN0_SERVICE" 127.0.0.1 $SOCKS_PROXY_PORT 2>> output.log ||
     true
      sudo -n networksetup -setsocksfirewallproxystate "$EN0_SERVICE" on 2>> output.log || true
    fi

    # IMPORTANT: Exclude local LAN and mDNS traffic from the proxy so P2P works natively
    # without source-IP masking from the Colima VM bridge.
    sudo -n networksetup -setproxybypassdomains "$svc" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
     172.16.0.0/12 "*.local" 2>> output.log || true
    if [ "$svc" != "$EN0_SERVICE" ]; then
      sudo -n networksetup -setproxybypassdomains "$EN0_SERVICE" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
     172.16.0.0/12 "*.local" 2>> output.log || true
    fi

    echo "done."
    echo "  All TCP traffic now routed through gateway -> WARP tunnel -> internet."
    echo "  (CLI tools: export ALL_PROXY=socks5://127.0.0.1:$SOCKS_PROXY_PORT)"
  else
    local svc="$1"
    echo "  (Linux global SOCKS proxy configuration is desktop-environment specific, skipping.)"
    echo "  SOCKS5 proxy is active and listening on 127.0.0.1:$SOCKS_PROXY_PORT"
  fi
}

disable_socks_proxy() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    for svc in "$ACTIVE_SERVICE" "$EN0_SERVICE"; do
      sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
    done
    echo "  SOCKS5 proxy disabled."
  else
    local svc="$1"
    echo "  (Linux global SOCKS proxy configuration skipped.)"
  fi
}

# Helper to restore host DNS to a clean state (used on script start, on failure,
# and after a successful STOP). Without this, a successful toggle-off leaves the
# host's resolver pinned to a now-dead gateway IP  every lookup stalls for macOS's
# DNS timeout (~5s) before falling through to whatever next server is configured.
# With it, we always leave the host in `1.1.1.1 / no search domain`.
#
# We set DNS on BOTH the active service AND the en0 service because macOS's scutil DNS
# resolver is commonly scoped to en0 (Wi-Fi). A per-service change on e.g.
# "USB 10/100 LAN" is ignored by the system resolver unless the en0 service is also updated.
reset_dns() {
  if [[ "$OSTYPE" == "darwin"* ]]; then
    if [ -n "$ACTIVE_SERVICE" ]; then
```

```
      for dns_svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
        networksetup -setsearchdomains "$dns_svc" "Empty" 2>> output.log || true
        networksetup -setdnsservers "$dns_svc" 1.1.1.1 2>> output.log || true
      done
    fi
  else
    if [ -n "$ACTIVE_SERVICE" ]; then
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      resolvectl revert "$ACTIVE_SERVICE" 2>/dev/null || true
    fi
  fi
}

# Verify basic direct-internet connectivity after a recovery/teardown: DNS via
# 1.1.1.1, then HTTP (curl) or ping to 1.1.1.1. Sets the global INTERNET_OK to
# "true" on any successful check (callers read $INTERNET_OK in their summary).
# Always returns 0 so a bare call can't trip `set -e`. Used by recover.sh's macOS
# and Linux recovery paths (extracted from a formerly-duplicated block).
verify_internet_connectivity() {
  step "Verifying internet connectivity"

  # Wait a moment for the network to settle
  sleep 2

  INTERNET_OK=false

  # Try DNS resolution first
  if host -W 3 google.com 1.1.1.1 >> output.log 2>&1; then
    ok "DNS resolution works (google.com via 1.1.1.1)"
    INTERNET_OK=true
  elif nslookup google.com 1.1.1.1 >> output.log 2>&1; then
    ok "DNS resolution works (nslookup google.com via 1.1.1.1)"
    INTERNET_OK=true
  else
    warn "DNS resolution check failed  will still try ping/curl"
  fi

  # Try actual HTTP connectivity
  if command -v curl >> output.log 2>&1; then
    if curl -sf --max-time 5 https://1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via HTTP"
      INTERNET_OK=true
    elif curl -sf --max-time 5 http://1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via HTTP (plain)"
      INTERNET_OK=true
    else
      warn "HTTP check failed"
    fi
  elif command -v ping >> output.log 2>&1; then
    if ping -c 1 -W 3 1.1.1.1 >> output.log 2>&1; then
      ok "Internet reachable via ping"
      INTERNET_OK=true
    else
      warn "Ping check failed"
    fi
  fi

  return 0
}

#  Local DNS Proxy (Python)
# When the tailnet data plane cannot establish, we use a Python DNS proxy.
# It listens on UDP:53, receives DNS queries, forwards them over TCP to
# AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (Docker port mapping), and returns the response.
# The Python script handles the DNS-over-TCP wire format (2-byte length prefix)
# correctly  unlike socat which just forwards raw bytes.
#
# Python3 is built into macOS  no external dependencies.
DNS_PROXY_PID=""

# Kill any leftover DNS proxy process
stop_local_dns_proxy() {
  if [ -n "$DNS_PROXY_PID" ]; then
    kill "$DNS_PROXY_PID" 2>/dev/null || true
    wait "$DNS_PROXY_PID" 2>/dev/null || true
    DNS_PROXY_PID=""
  fi
  # Clean up any stale Python DNS proxy processes
  sudo -n pkill -f "dns-proxy.py" 2>/dev/null || true
}

```

```
# Start local DNS proxy (Python)
start_local_dns_proxy() {
  if [ -z "$DNS_PROXY_BIN" ]; then
    echo "  Python not found. Python3 is built into macOS  this shouldn't happen."
    return 1
  fi

  if [ ! -f "$DNS_PROXY_SCRIPT" ]; then
    echo "  DNS proxy script not found at $DNS_PROXY_SCRIPT"
    return 1
  fi

  # Kill any leftover process first
  stop_local_dns_proxy

  echo -n "  Starting local DNS proxy via Python (UDP:53  TCP:${DNS_PROXY_PORT})... "
  sudo -n bash -c "TARGET_PORT=$DNS_PROXY_PORT \"$DNS_PROXY_BIN\" \"$DNS_PROXY_SCRIPT\"" &
  DNS_PROXY_PID=$!
  disown "$DNS_PROXY_PID" 2>/dev/null || true

  # Give it a moment to bind
  sleep 0.5

  if kill -0 "$DNS_PROXY_PID" 2>/dev/null; then
    echo "started (PID $DNS_PROXY_PID)."

    # Hijack host DNS to localhost
    echo -n "  Hijacking host DNS to 127.0.0.1 for ad-blocking... "
    if [[ "$OSTYPE" == "darwin"* ]]; then
      networksetup -setsearchdomains "$ACTIVE_SERVICE" "Empty" 2>/dev/null || true
      networksetup -setsearchdomains "$ENO_SERVICE" "Empty" 2>/dev/null || true
      if ! networksetup -setdnsservers "$ACTIVE_SERVICE" 127.0.0.1; then
        echo "FAILED (networksetup error  check permissions)."
        stop_local_dns_proxy
        return 1
      fi
      networksetup -setdnsservers "$ENO_SERVICE" 127.0.0.1 2>/dev/null || true
      sudo -n dscacheutil -flushcache 2>/dev/null || true
      sudo -n killall -HUP mDNSResponder 2>/dev/null || true
    else
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
      if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" 127.0.0.1; then
        echo "FAILED (resolvectl error  check permissions)."
        stop_local_dns_proxy
        return 1
      fi
      resolvectl dns "$ACTIVE_SERVICE" 127.0.0.1 2>/dev/null || true
      sudo -n resolvectl flush-caches 2>/dev/null || true
    fi

    # Verify DNS actually works through the proxy
    echo -n "  Verifying DNS resolution... "
    if dig @127.0.0.1 google.com +short +timeout=5 &>/dev/null; then
      echo "ok."
      return 0
    else
      echo "FAILED (DNS queries not reaching AdGuard)."
      echo "  Restoring DNS to 1.1.1.1..."
      reset_dns
      stop_local_dns_proxy
      return 1
    fi
  else
    echo "FAILED (could not bind port 53  is it in use?)."
    DNS_PROXY_PID=""
    return 1
  fi
}

# Helper to FORCIBLY hijack host DNS to a single server (the gateway's AdGuard IP)
# and VERIFY via `networksetup -getdnsservers` that the only name server is exactly
# $1. Returns 0 iff the live state matches; non-zero if it can't be confirmed.
#
# This deliberately does NOT append a 1.1.1.1 fallback: macOS's resolver queries
# the list in order and falls back to the next entry on timeout, so on a
# misbehaving AdGuard (filter compile crash, OOM kill, etc.) the resolver still
# leaks to 1.1.1.1 and silently bypasses ad-blocking. With no fallback, a broken
# gateway manifests as a *visible* DNS outage that prompts the user to
# investigate  which is the correct failure mode.
force_dns_to_gateway() {
```

```
  if [[ "$OSTYPE" == "darwin"* ]]; then
    local target_ip="$1"
    if [ -z "$target_ip" ] || [ -z "$ACTIVE_SERVICE" ]; then
      echo "  [force_dns] ERROR: target_ip or ACTIVE_SERVICE is empty" >&2
      return 1
    fi

    local attempt
    for attempt in 1 2 3; do
      echo -n "  [force_dns] Attempt $attempt/3: setting DNS to $target_ip on \"$ACTIVE_SERVICE\" + \"$ENO_SERVICE\"... "

      # Apply to BOTH the active service AND the en0 service  the scutil DNS resolver
      # is scoped to en0 (Wi-Fi) and ignores per-service changes that don't include it.
      # Fail fast on ACTIVE_SERVICE; best-effort on ENO_SERVICE (it may not exist).
      networksetup -setsearchdomains "$ACTIVE_SERVICE" "ts.net" || true
      if ! networksetup -setdnsservers "$ACTIVE_SERVICE" "$target_ip"; then
        echo "FAILED (networksetup error  check permissions)"
        sleep 1
        continue
      fi
      networksetup -setsearchdomains "$ENO_SERVICE" "ts.net" || true
      networksetup -setdnsservers "$ENO_SERVICE" "$target_ip" || true

      local entries
      entries=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log \
                | awk '/^(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)$/ {print}')
      if echo "$entries" | grep -qFx "$target_ip"; then
        echo "VERIFIED"
        return 0
      fi

      local current
      current=$(echo "$entries" | tr '\n' ' ' | sed 's/ $//')
      echo "NOT YET (current DNS: [${current:-none}])"

      if [ "$attempt" = "3" ]; then sleep 2; else sleep 1; fi
    done

    echo "  [force_dns] FAILED after 3 attempts"
    return 1
  else
    local target_ip="$1"

    for attempt in {1..3}; do
      echo -n "  [force_dns] Attempt $attempt/3: setting DNS to $target_ip on $ACTIVE_SERVICE... "
      sudo -n resolvectl domain "$ACTIVE_SERVICE" "~ts.net" 2>/dev/null || true
      if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" "$target_ip"; then
        echo "FAILED (resolvectl error)"
      else
        # Verify
        entries=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+')
        if [ "$entries" == "$target_ip" ]; then
          echo "VERIFIED"
          return 0
        else
          echo "MISMATCH (Resolver refused lock)"
        fi
      fi
      sleep 2
    done
    return 1
  fi
}
```

12 scripts/watcher.sh

```
#!/bin/bash
# scripts/watcher.sh  nullexit post-wake / post-roam recovery watcher
#
# Long-running daemon. Launched by launchd as a LaunchAgent at
# ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
# (label com.nullexit.wake-recovery).
#
# Two independent listeners run as backgrounded pipelines:
```

```
#
#   (1) WAKE listener  `log stream` live-follows the macOS unified log for
#       power-management events (lid closewake, manual wake, DarkWake from
#       clamshell). Each event triggers run_recover (with a global debounce so
#       a single wake that emits 2-3 log lines only causes ONE post-wake run).
#
#   (2) NETWORK listener  `scutil n.watch` on `State:/Network/Global/IPv4`
#       fires on every Wi-Fi roam, ethernet swap, hotspot join, captive-portal
#       re-bind, VPN up/down, etc. Same debounce.
#
# Both listeners call run_recover, which:
#   - checks /tmp/nullexit-gateway-active.marker (only act when toggle.sh left
#     the gateway up)
#   - debounces (10s default) so multiple events don't pile up
#   - shells out to `bash <repo>/recover.sh --post-wake` directly (no GUI auth prompt needed due to NOPASSWD
    sudoers)
#
# See devref.md 10.29 for the full Apple power management primer and the
# rationale behind using `log stream` + `scutil n.watch` from a shell daemon.

# NOTE: errexit (`set -e`) is intentionally NOT enabled. This is a long-running
# daemon, not a discrete-step script, and errexit actively breaks the reconnect
# loops below: any failed pipe (log stream on rotation, scutil on state churn)
# would kill the outer subshell before `echo "reconnecting..."; sleep 5;` runs,
# AND the parent's final `wait` would propagate a non-zero child exit and kill
# the entire watcher. Every error path is already wrapped in `|| true` /
# `2>/dev/null` fallbacks.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
source "$SCRIPT_DIR/common.sh"
setup_standard_path

LOG="$SCRIPT_DIR/../output.log"
# Resolve our own directory; recover.sh lives one level up at the repo root.
# This makes the daemon install-agnostic  no need to template the path
# in launchd, no need to keep a deploy-specific hardcoded location in sync.
RECOVER="$SCRIPT_DIR/../recover.sh"
MARKER="$MARKER_FILE"
DEBOUNCE_FILE="/tmp/nullexit-watcher.last-recovery"
DEBOUNCE_SECONDS="${NULLEXIT_DEBOUNCE_SECONDS:-10}"

exec >> "$LOG" 2>&1

#  Single-instance lock
# Without this, repeated lid-closewake cycles under macOS Sonoma+ accumulate
# orphan watcher.sh processes: launchd suspends on sleep and SIGCONT-resumes
# on wake, and KeepAlive.Crashed=true caused relaunch on any non-zero exit;
# the listener grandchildren (log stream | while ...  and  (echo n.add ...;
# sleep 86400) | scutil | while ...) stay alive because pipe producers don't
# reliably respond to plain SIGTERM if the consumer is in a half-closed state.
#
# We solve it with a PID-file lock (POSIX-portable; macOS does NOT ship `flock`).
# The trap below removes the lock file on every exit so a stale lock never
# outlives a crashed watcher. On launch: read the existing PID, check it is
# alive AND is actually a `scripts/watcher.sh` process; if so, exit 0; else
# overwrite with our PID and proceed. The check-then-write window has a TOCTOU
# race but is microseconds wide and only matters for hand-launched duplicate
# test invocations; production is single-launch via launchctl load -w.
LOCK_FILE="/tmp/nullexit-watcher.lock"
LOCK_PID=""
if [ -e "$LOCK_FILE" ]; then
  LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
  if [ -n "$LOCK_PID" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
    # Verify the holder is actually a scripts/watcher.sh process (not some
    # unrelated process that got the recycled PID). The exact-launchd pattern
    # `bash <abs-path>/scripts/watcher.sh` is hard to accidentally match with
    # a `grep scripts/watcher.sh .` invocation because we anchor on `^bash`
    # followed by whitespace and end-of-line on the script path.
    if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -qE "(^|[[:space:]])bash[[:space:]]+.*scripts/watcher\.sh$
     "; then
      echo "[$(date -u +%FT%TZ)] watcher.sh: another instance holds $LOCK_FILE (pid=$LOCK_PID), exiting"
      exit 0
    fi
  fi
fi
echo $$ > "$LOCK_FILE"

echo "[$(date -u +%FT%TZ)] watcher.sh start (pid=$$ ppid=$PPID user=$(id -un) lock=acquired)"

# Treat launchd-resume as a wake signal. Apple's launchd SUSPENDS LaunchAgents
# during system sleep and restores them on wake; during that suspended gap,
# macOS's `log stream` cannot catch the wake event that prompted the resume.
```

```
# The next thing that happens after wake IS our own startup, so we always
# fire one post-wake run unconditionally here. The marker check + debounce
# in run_recover() make this idempotent: if the gateway isn't up or we just
# debounced, it no-ops. Without this, the FIRST wake-after-install goes
# unseen and the watcher appears broken until the SECOND wake.
#  Helper: run recover.sh --post-wake if the gateway is active
run_recover() {
  local why="$1"
  if [ ! -f "$MARKER" ]; then
    echo "[$(date -u +%FT%TZ)] $why  no $MARKER, skip"
    return 0
  fi
  local now last
  now=$(date +%s)
  last=$(cat "$DEBOUNCE_FILE" 2>/dev/null || echo 0)
  if [ "$((now - last))" -lt "$DEBOUNCE_SECONDS" ]; then
    echo "[$(date -u +%FT%TZ)] $why  debounced ($((now - last))s since last, threshold ${DEBOUNCE_SECONDS}s)"
    return 0
  fi
  echo "$now" > "$DEBOUNCE_FILE"
  echo "[$(date -u +%FT%TZ)] $why  bash $RECOVER --post-wake"
  bash "$RECOVER" --post-wake
  local exit_code=$?
  echo "$(date +%s)" > "$DEBOUNCE_FILE"
  echo "[$(date -u +%FT%TZ)] $why  exit=$exit_code (now=$(date +%s))"
  return $exit_code
}

if [ -f "$MARKER" ]; then
  run_recover "initial post-wake (launchd resume = wake signal)"
fi

#  LAN P2P Auto-Detection
# Detects whether the current network allows safe Tailscale LAN P2P by:
#   1. Checking WPA2-Enterprise (802.1x) via system_profiler SPAirPortDataType  always forces P2P=false
#      because enterprise networks almost universally have AP Client Isolation.
#   2. AP Isolation probe  sends a broadcast ping and checks if ANY LAN host responds.
#      If no LAN hosts respond on a non-empty network, AP Isolation is active.
# Writes 'true' or 'false' to .lan_p2p_detected in the repo root.
# routing-fix.sh reads this file dynamically every 30s loop iteration.
P2P_OVERRIDE_FILE="$SCRIPT_DIR/../.lan_p2p_detected"


detect_lan_p2p_mode() {
  local reason="unknown" allow="false"

  # Step 1: Detect Wi-Fi security type via system_profiler (airport binary removed in macOS 14.4+)
  # SPAirPortDataType lists all known networks; the first "Security:" line is the current association.
  local security
  security=$(system_profiler SPAirPortDataType 2>/dev/null \
    | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF; exit}')

  if [ -z "$security" ]; then
    # Not on Wi-Fi or system_profiler unavailable  fail safe
    reason="no-wifi" allow="false"
  elif echo "$security" | grep -qi "Enterprise"; then
    # WPA2/WPA3 Enterprise = 802.1x  always AP isolated
    reason="802.1x-enterprise" allow="false"
  elif echo "$security" | grep -qi "Personal\|None\|Open"; then
    # WPA2/WPA3 Personal or open  probe for AP isolation
    local gw
    gw=$(route -n get 1.1.1.1 2>/dev/null | awk '/gateway:/{print $2}')
    if [ -z "$gw" ]; then
      reason="no-default-route" allow="false"
    else
      local responses
      responses=$(ping -c 3 -t 1 -q "$gw" 2>/dev/null | awk '/packets transmitted/{print $4}')
      if [ "${responses:-0}" -gt 0 ]; then
        reason="wpa-personal-lan-reachable" allow="true"
      else
        reason="wpa-personal-ap-isolated" allow="false"
      fi
    fi
  else
    # Unknown security type  fail safe
    reason="unknown-security-type" allow="false"
  fi

  echo "[$(date -u +%FT%TZ)] P2P detect: security='${security:-none}'  allow=${allow} (reason: ${reason})"
  echo "$allow" > "$P2P_OVERRIDE_FILE"
  write_host_ips
```

```
}

# Run once on startup
detect_lan_p2p_mode

#  Listener 1: WAKE events via unified log
# Predicate picks up:
#   * "didWake"              regular kernel wake events (lid open, RTC alarm)
#   * "Wake from Sleep"      powerd-driven normal wake
#   * "DarkWake"             clamshell-display wake that didn't fully wake the
#                             Mac (relevant when lid opens during display sleep)
#   * "Waking from"          generic catch-all for variations in newer macOS
# `[c]` makes the match case-insensitive (the unified log uses mixed-case
# phrases). --info reduces noise by dropping debug/trace log levels.
# Subsystem-based predicate catches every powermanagement emission (willSleep,
# didWake, didDim, didUndim, DarkWake, etc.) regardless of message phrasing
# across kernel versions. Phrasing-based predicates ('didWake', 'Wake from Sleep')
# are fragile across macOS releases. We accept the extra log noise because the
# run_recover() debounce + marker check filter out anything that isn't a real
# gateway-relevant event.
(
  # Reconnect loop: macOS rotates the unified log buffer periodically; when
  # that happens, this `log stream` subscriber's pipe EOFs, the inner
  # `while read` consumer terminates, and without this wrapper the listener
  # would be silently dead for the rest of the watcher's lifetime.
  while :; do
    log stream --predicate \
      'subsystem == "com.apple.powermanagement"' \
       --style compact --info 2>/dev/null \
      | while IFS= read -r line; do
          [ -z "$line" ] && continue
          run_recover "WAKE: $(echo "$line" | head -c 100)"
        done
    echo "[$(date -u +%FT%TZ)] log stream subshell exited; reconnecting in 5s"
    sleep 5
  done
) &

#  Listener 2: NETWORK state changes via scutil
# `scutil n.watch` on a state key prints SCEventUpdate lines whenever the key
# changes. We watch Global/IPv4 because that's the umbrella that catches link,
# address, route, and DNS changes for ALL interfaces in one namespace.
# The pipe is kept alive by an `sleep 86400` loop so scutil never sees EOF
# from us (which would silently terminate the watch).
(
  # Reconnect loop: scutil can exit noisily on certain state changes (rare,
  # but observed). Without the wrapper, the network listener's pipe EOFs and
  # every subsequent roam goes unseen until launchd restarts us.
  while :; do
    (
      # Watch both IPv4 and IPv6 umbrella state  IPv6-only or dual-stack
      # networks never fire on the IPv4 key alone, and the user's first roam
      # in a hotel / airport / on cellular hotspot is often IPv6-first under
      # NAT64.
      echo "n.add State:/Network/Global/IPv4"
      echo "n.add State:/Network/Global/IPv6"
      echo "n.watch"
      while true; do sleep 86400; done
    ) | script -q /dev/null scutil 2>/dev/null \
      | while IFS= read -r line; do
          case "$line" in
            *changedKey*|*State:/Network/Global/IPv*|*n.state*|*SCEventUpdate*)
              detect_lan_p2p_mode
              run_recover "NET: $(echo "$line" | head -c 100)"
              ;;
            *) ;;
          esac
        done
    echo "[$(date -u +%FT%TZ)] scutil pipe exited; reconnecting in 5s"
    sleep 5
  done
) &
#  Listener 3: TUNNEL_FAILED_CLOSED.marker Watcher
# Polls for the fail-closed marker dropped by the routing-fix container.
(
  last_notified=0
  while :; do
    if [ -f "$SCRIPT_DIR/../TUNNEL_FAILED_CLOSED.marker" ]; then
      now=$(date +%s)
      # Notify once every 5 minutes (300s) if it stays failed
      if [ "$((now - last_notified))" -ge 300 ]; then
```

```
        osascript -e 'display notification "VPN Tunnel Health Check Failed. Internet traffic is blackholed to prevent IP leak." with title "NullExit Alert"' || true
        last_notified=$now
      fi
    else
      last_notified=0
    fi
    sleep 3
  done
) &

#  Lifetime
# `wait` blocks until both background pipelines exit (which they normally
# never do). launchctl `bootout`/SIGTERM triggers the trap, which kills the
# whole process group so the sleep-forever in the scutil pipe also dies.
# We catch TERM/INT/HUP for explicit signals and EXIT for any normal-exit path
# (so listeners always get cleaned up even if `exit 0` runs somewhere).
# `rm -f $LOCK_FILE` first so a stale lock can never block the next launch.
# `pkill -TERM -P $$` then targets direct listener-children, then `kill 0`
# takes care of any backgrounded group-mates not reachable via the parent.
trap 'echo "[$(date -u +%FT%TZ)] watcher.sh terminating (signal/exit=$?)"; rm -f "$LOCK_FILE" 2>/dev/null || true; pkill -TERM -P $$ 2>/dev/null || true; kill 0 2>/dev/null || true; exit 0' TERM INT HUP EXIT
wait
```

13 scripts/crypto.sh

```
#!/usr/bin/env bash
# crypto.sh - Cryptographic Integrity Check

set -euo pipefail
cd "$(dirname "$0")/.."

if [ "$#" -ne 1 ] || { [ "$1" != "--sign" ] && [ "$1" != "--verify" ]; }; then
  echo "Usage: $0 --sign | --verify"
  exit 1
fi

if [ ! -f .env ]; then
  echo "Error: .env not found."
  exit 1
fi

SEED=$(grep "^NULLEXIT_SEED=" .env | cut -d '=' -f2)
if [ -z "$SEED" ]; then
  echo "Error: NULLEXIT_SEED not found in .env."
  exit 1
fi

if [ "$1" == "--sign" ]; then
  echo "Generating cryptographic signatures using seed..."
  rm -f .signatures
  for file in toggle.sh recover.sh scripts/common.sh scripts/routing-fix.sh scripts/diagnose-host-leak.sh scripts/watcher.sh scripts/pf.conf post-rules.txt docker-compose.yml black_list.txt white_list.txt; do
    if [ -f "$file" ]; then
      hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
      echo "$file:$hash" >> .signatures
      echo "  [Signed] $file"
    else
      echo "  [Skipped] $file not found"
    fi
  done
  echo "Signatures saved to .signatures"
  exit 0
fi

if [ "$1" == "--verify" ]; then
  if [ ! -f .signatures ]; then
    echo "Error: .signatures file missing. Run $0 --sign first."
    exit 1
  fi
  FAIL=0
  while IFS=: read -r file expected_hash; do
    if [ -z "$file" ]; then continue; fi
    if [ ! -f "$file" ]; then
      echo "[FAIL] $file is missing!"
      FAIL=1
```

```
      continue
    fi
    actual_hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
    if [ "$actual_hash" != "$expected_hash" ]; then
      echo "[FAIL] $file has been modified! Hash mismatch."
      FAIL=1
    fi
  done < .signatures

  if [ "$FAIL" -eq 1 ]; then
    echo ""
    echo "CRITICAL: Script integrity verification failed!"
    echo "One or more core files have been modified without authorization."
    echo "To authorize these changes, run: ./scripts/crypto.sh --sign"
    echo ""
    exit 1
  fi
  exit 0
fi
```

14 .aiignore

```
# Environment variables and secrets
.env
tor-bridges.txt
```

15 .claude/settings.local.json

```
{
  "permissions": {
    "allow": [
      "Bash(grep *)"
    ]
  }
}
```

16 .gitignore

```
# Environment variables and secrets
.env

# wgcf generated files containing keys and account info
wgcf-account.toml
wgcf-profile.conf

# Tailscale persistent state containing node identity
tailscale-state/

# Binaries
wgcf

# macOS App bundles
*.app/
*.app

# AdGuard Home persistent state
adguard/
.gateway_ip
logs.txt

#
stability_incident_report.md
output.log
.DS_Store
*.desktop
blocked.log
wtf.md
host-leak-diagnostic-*.txt
```

```
nullexit_review.md
nullexit_unified.tex
nullexit_unified.txt

nullexit_unified.synctex.gz
tor-bridges.txt

# Runtime state files
.host_ips
.lan_p2p_detected
.build_hash
.rules_hash
.signatures
TUNNEL_FAILED_CLOSED.marker

# Python cache
__pycache__/
*.py[cod]

# LaTeX build artifacts
*.aux
*.log
*.out
*.toc
*.xdv
*.fdb_latexmk
*.fls
*.synctex*
refactor.md
```

17 Recover-Gateway.applescript

```
tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
try
  do shell script "true" with administrator privileges
on error
  return
end try
tell application "Terminal"
  activate
  do script "cd " & quoted form of scriptDir & "; bash recover.sh"
end tell
```

18 Toggle-Gateway.applescript

```
tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
try
  do shell script "true" with administrator privileges
on error
  return
end try
tell application "Terminal"
  activate
  do script "cd " & quoted form of scriptDir & "; ./toggle.sh"
end tell
```

19 benchmarks/benchmark.sh

```
#!/usr/bin/env bash

# benchmark.sh
# Automates the full benchmarking suite (both Python download test and Lighthouse HTML test)
# Runs the suite with Gateway ON, toggles OFF, runs again, then toggles back ON.

set -e

# Change to project root so we can access toggle.sh reliably
cd "$(dirname "$0")/.."
```

```
11
12 # Ensure jq is installed
13 if ! command -v jq &> /dev/null; then
14     echo "Error: 'jq' is not installed. Please run 'brew install jq'"
15     exit 1
16 fi
17
18 run_lighthouse() {
19     local url=$1
20     local file_prefix=$2
21     echo "  [Lighthouse] Testing $url ..."
22
23     # Run lighthouse headlessly, silencing standard output
24     npx -y lighthouse "$url" --chrome-flags="--headless" --only-categories=performance --output=json --output-
        path="${file_prefix}.json" >/dev/null 2>&1
25
26     # Parse the results
27     local score=$(jq -r '.categories.performance.score' "${file_prefix}.json")
28     local tti=$(jq -r '.audits["interactive"].displayValue' "${file_prefix}.json")
29     local si=$(jq -r '.audits["speed-index"].displayValue' "${file_prefix}.json")
30
31     # Convert score to percentage
32     local score_pct=$(awk "BEGIN {print $score*100}")
33
34     echo "    -> Performance Score: ${score_pct}%"
35     echo "    -> Time to Interactive: $tti"
36     echo "    -> Speed Index: $si"
37
38     # Cleanup
39     rm -f "${file_prefix}.json"
40 }
41
42 echo "============================================="
43 echo "   PHASE 1: Testing with Gateway ON"
44 echo "============================================="
45 # 1. Run raw Python download test
46 python3 benchmarks/test_load.py
47
48 # 2. Run Lighthouse HTML render test
49 echo ""
50 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
51 run_lighthouse "https://www.cnet.com/" "on_cnet"
52 run_lighthouse "https://www.independent.co.uk/" "on_ind"
53
54
55 echo ""
56 echo "============================================="
57 echo "   PHASE 2: Toggling Gateway OFF"
58 echo "============================================="
59 ./toggle.sh
60 echo "Waiting 10 seconds for macOS Wi-Fi to stabilize after DNS reset..."
61 sleep 10
62
63
64 echo ""
65 echo "============================================="
66 echo "   PHASE 3: Testing with Gateway OFF"
67 echo "============================================="
68 # 1. Run raw Python download test
69 python3 benchmarks/test_load.py
70
71 # 2. Run Lighthouse HTML render test
72 echo ""
73 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
74 run_lighthouse "https://www.cnet.com/" "off_cnet"
75 run_lighthouse "https://www.independent.co.uk/" "off_ind"
76
77
78 echo ""
79 echo "============================================="
80 echo "   PHASE 4: Restoring Gateway ON"
81 echo "============================================="
82 ./toggle.sh
83
84 echo ""
85 echo "============================================="
86 echo " BENCHMARKS COMPLETE!"
87 echo "You can scroll up to compare Phase 1 (AdGuard ON) vs Phase 3 (AdGuard OFF)."
88 echo "============================================="
```

20 benchmarks/test_load.py

```python
import urllib.request
import urllib.error
import time
from html.parser import HTMLParser
import concurrent.futures

class ResourceParser(HTMLParser):
    def __init__(self):
        super().__init__()
        self.urls = set()

    def handle_starttag(self, tag, attrs):
        attrs = dict(attrs)
        url = None
        if tag in ['img', 'script']:
            url = attrs.get('src')
        elif tag == 'link' and attrs.get('rel') == 'stylesheet':
            url = attrs.get('href')

        if url and url.startswith('http'):
            self.urls.add(url)

def fetch_url(url):
    try:
        req = urllib.request.Request(url, headers={'User-Agent': 'Mozilla/5.0'})
        urllib.request.urlopen(req, timeout=3)
        return True
    except Exception:
        return False

def measure_url(target_url):
    print(f"\nTesting {target_url} Load...")
    start_total = time.time()

    # Fetch main HTML
    try:
        req = urllib.request.Request(target_url, headers={'User-Agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7) AppleWebKit/537.36'})
        response = urllib.request.urlopen(req, timeout=5)
        html = response.read().decode('utf-8', errors='ignore')
    except Exception as e:
        print(f"Failed to fetch {target_url}: {e}")
        return

    parser = ResourceParser()
    parser.feed(html)
    urls = list(parser.urls)
    print(f"Found {len(urls)} subresources to fetch (scripts, images, css)...")

    # Concurrently fetch subresources
    success_count = 0
    fail_count = 0
    max_threads = min(64, max(1, len(urls)))
    with concurrent.futures.ThreadPoolExecutor(max_workers=max_threads) as executor:
        results = list(executor.map(fetch_url, urls))

    success_count = sum(1 for r in results if r)
    fail_count = len(results) - success_count

    total_time = time.time() - start_total
    print(f"Time taken: {total_time:.2f} seconds")
    print(f"Successfully fetched: {success_count}, Failed/Blocked: {fail_count}")

urls_to_test = [
    "https://www.dailymail.co.uk/home/index.html",
    "https://www.cnet.com/",
    "https://www.independent.co.uk/"
]

for u in urls_to_test:
    measure_url(u)
```

21 black_list.txt

```
ad.doubleclick.net
adclick.g.doubleclick.net
adeventtracker.spotify.com
adfstat.yandex.ru
ads-api.tiktok.com
ads-api.twitter.com
ads-fa.spotify.com
ads-sg.tiktok.com
ads.tiktok.com
adserver.unityads.unity3d.com
adsfs.oppomobile.com
an.facebook.com$important
connect.facebook.net$important
pixel.facebook.com$important
analytics.query.yahoo.com
analytics.s3.amazonaws.com
analytics.spotify.com
analyticsengine.s3.amazonaws.com
appmetrica.yandex.ru
audio2.spotify.com
bdapi-ads.realmemobile.com
bdapi-in-ads.realmemobile.com
bounceexchange.com
bs.serving-sys.com
business-api.tiktok.com
ck.ads.oppomobile.com
crashdump.spotify.com
data.ads.oppomobile.com
data.mistat.india.xiaomi.com
data.mistat.rus.xiaomi.com
data.mistat.xiaomi.com
desktop.spotify.com
doubleclick.net
ds.metric.spotify.com
gemini.yahoo.com
googleads.g.doubleclick.net
googleadservices.com
googlesyndication.com
grs.hicloud.com
gtm.spotify.com
iot-eu-logser.realme.com
iot-logser.realme.com
log.byteoversea.com
log.fc.yahoo.com
log.spotify.com
m.doubleclick.net
mediavisor.doubleclick.net
metrics.spotify.com
metrika.yandex.ru
omaze.com
pagead2.googlesyndication.com
securepubads.g.doubleclick.net
static.doubleclick.net
stats.g.doubleclick.net
tracking.rus.miui.com
udcm.yahoo.com
v.jwpcdn.com
weblb-wg.gslb.spotify.com
webview.unityads.unity3d.com
diagassets.apple.com
iphonesubmissions.apple.com
radarsubmissions.apple.com
metrics.apple.com
xp.apple.com
heads-fa.spotify.com
heads4.spotify.com
pixel.spotify.com
segs.spotify.com
audio-fa.spotify.com
events.data.microsoft.com
datarouter-prod.ak.epicgames.com
nel.cloudflare.com
eu-mobile.events.data.microsoft.com
stocks-analytics-events.apple.com
```

22 cloud-init.yaml

```
#cloud-config
# Generic Cloud-Init Bootstrap Script for Remote Exit Nodes (Ubuntu/Debian)
# Paste this into the "User Data" or "Initialization Script" section when creating a VPS.

package_update: true
package_upgrade: true

packages:
  - apt-transport-https
  - ca-certificates
  - curl
  - gnupg
  - lsb-release
  - git
  - python3
  - python3-pip

runcmd:
  # 1. Install Docker natively
  - curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg
  - echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
  - sudo apt-get update
  - sudo apt-get install -y docker-ce docker-ce-cli containerd.io docker-compose-plugin

  # 2. Add default user (ubuntu) to docker group
  - usermod -aG docker ubuntu || true
  - usermod -aG docker debian || true

  # 3. Create the installation directory
  - mkdir -p /opt/nullexit
  - chown -R ubuntu:ubuntu /opt/nullexit || chown -R debian:debian /opt/nullexit

  # 4. Enable IPv4 and IPv6 forwarding for VPN routing
  - echo "net.ipv4.ip_forward=1" >> /etc/sysctl.d/99-tailscale.conf
  - echo "net.ipv6.conf.all.forwarding=1" >> /etc/sysctl.d/99-tailscale.conf
  - sysctl -p /etc/sysctl.d/99-tailscale.conf

final_message: "The remote exit node has been fully bootstrapped. SSH in, clone the repo to /opt/nullexit, add your .env, and run ./setup-linux.sh!"
```

23 docker/tor/Dockerfile

```
FROM debian:bookworm-slim

RUN apt-get update && \
    apt-get install -y tor obfs4proxy gosu bash && \
    rm -rf /var/lib/apt/lists/*

COPY entrypoint.sh /entrypoint.sh
RUN chmod +x /entrypoint.sh

ENTRYPOINT ["/entrypoint.sh"]
```

24 docker/tor/entrypoint.sh

```
#!/bin/bash
set -e

TORRC="/etc/tor/torrc"
mkdir -p /etc/tor /var/lib/tor

# Base configuration: SOCKS 9050, Control 9051, TransPort 9040, DNSPort 5353  all bound to
# 0.0.0.0 (required for Docker/Colima host port-mapping). The ControlPort is secured via a
# dynamic HashedControlPassword (added below when TOR_PASSWORD is set); see devref.md 15.10.2.
cat <<EOF > $TORRC
SocksPort 0.0.0.0:9050
```

```
ControlPort 0.0.0.0:9051
CookieAuthentication 0
Log notice stdout
DataDirectory /var/lib/tor

# Transparent proxying & DNS resolution for .onion mapping
TransPort 0.0.0.0:9040
DNSPort 0.0.0.0:5353
AutomapHostsOnResolve 1
VirtualAddrNetworkIPv4 198.18.0.0/15
EOF

if [ -n "$TOR_PASSWORD" ]; then
    HASH=$(tor --hash-password "$TOR_PASSWORD" | tail -n 1)
    echo "HashedControlPassword $HASH" >> $TORRC
fi

if [ -n "$TOR_EXCLUDE_EXIT_NODES" ]; then
    echo "ExcludeExitNodes $TOR_EXCLUDE_EXIT_NODES" >> $TORRC
    echo "StrictNodes 1" >> $TORRC
fi

if [ "${TOR_USE_BRIDGES:-false}" = "true" ]; then
    BRIDGE_FILE="${TOR_BRIDGE_LINES_FILE:-/tor-bridges.txt}"

    if ! grep -vE "^[[:space:]]*#" "$BRIDGE_FILE" | grep -q '[^[:space:]]'; then
        MSG="[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL] [Tor Proxy] TOR_USE_BRIDGES is enabled, but no bridge lines found at $BRIDGE_FILE. Proxy startup ABORTED. Get bridges from https://bridges.torproject.org."
        echo "$MSG"
        if [ -w "/output.log" ]; then
            echo "$MSG" >> /output.log
        fi
        # Sleep indefinitely to prevent docker-compose 'unless-stopped' from triggering a crash loop
        exec sleep infinity
    fi

    echo "UseBridges 1" >> $TORRC
    echo "ClientTransportPlugin obfs4 exec /usr/bin/obfs4proxy" >> $TORRC

    while IFS= read -r line; do
        # Ignore comments and empty lines
        [[ -z "$line" || "$line" == \#* ]] && continue
        echo "Bridge $line" >> $TORRC
    done < "$BRIDGE_FILE"
fi

# Debian tor package uses 'debian-tor' user
chown -R debian-tor:debian-tor /var/lib/tor /etc/tor

exec gosu debian-tor /usr/bin/tor -f $TORRC
```

25 launchd/com.nullexit.wake-recovery.plist

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
<plist version="1.0">
<dict>
    <key>Label</key>
    <string>com.nullexit.wake-recovery</string>

    <!-- WorkingDirectory + relative program arg keeps this plist install-agnostic.
         The source-of-truth for the install path is __NULLEXIT_HOME__, which
         setup.sh or the user must substitute with their absolute nullexit
         install root at install time. Program arg is `scripts/watcher.sh`
         RELATIVE to WorkingDirectory, so launchd's evaluated cwd IS the
         install root  and BASH_SOURCE inside scripts/watcher.sh still
         resolves to the real <install>/scripts/watcher.sh, preserving the
         SCRIPT_DIR pattern. After install-time `sed -i '' 's|__NULLEXIT_HOME__|<abs_path>|'`,
         the plist is portable to any clone location. -->
    <key>WorkingDirectory</key>
    <string>__NULLEXIT_HOME__</string>

    <key>ProgramArguments</key>
    <array>
        <string>/bin/bash</string>
        <string>scripts/watcher.sh</string>
```

```
    </array>

    <!-- Start the daemon immediately on launchctl load (instead of waiting
         for the next user login). Combined with keepAlive below this means
         once installed: load once, runs forever across logouts/reboots. -->
    <key>RunAtLoad</key>
    <true/>

    <!-- Restart policy:
         SuccessfulExit=false   SIGTERM / clean exit does NOT relaunch (so
                                `launchctl bootout` cleanly stops, not loops).
         Crashed=false         HARD crashes ALSO do NOT relaunch. We've seen
                                that repeated sleep/wake cycles on macOS
                                Sonoma+ accumulate orphan watcher.sh PIDs
                                because SIGCONT/resume doesn't reliably kill
                                the prior instance. The watcher.sh script
                                now enforces single-instance via a manual
                                PID-file and process-identity check at the
                                top, so a re-launch on crash is
                                actively dangerous (it would skip the lock
                                and exit; better to surface the crash in
                                logs than to silently 0-process). -->
    <key>KeepAlive</key>
    <dict>
        <key>SuccessfulExit</key>
        <false/>
        <key>Crashed</key>
        <false/>
    </dict>

    <!-- Background hint to App Nap so it doesn't get suspended when
         the user is inactive. We're the entire reason this LaunchAgent
         needs to NOT nap. -->
    <key>ProcessType</key>
    <string>Background</string>

    <!-- launchd-level stdout/stderr (the bash launcher shell). The script
         itself writes to /tmp/nullexit-watcher.log via exec >>. These
         separate channels help us tell "launchd crashed" from "the bash
         wrapper piped to recover.sh errored". -->
    <key>StandardOutPath</key>
    <string>/tmp/nullexit-watcher.out.log</string>
    <key>StandardErrorPath</key>
    <string>/tmp/nullexit-watcher.err.log</string>
</dict>
</plist>
```

26 post-rules.txt

```
iptables -A FORWARD -i tailscale0 -o tun0 -j ACCEPT
iptables -A FORWARD -i tun0 -o tailscale0 -j ACCEPT
iptables -A INPUT -i tailscale0 -j ACCEPT
iptables -A INPUT -i eth0 -p udp --dport 41642 -j ACCEPT
iptables -t nat -A POSTROUTING -o tun0 -j MASQUERADE
# Clamp MSS to 1120. With Tailscale overhead on top of WARP (each WireGuard encapsulation adds ~60 bytes),
# the safe MSS for double-tunneled traffic is 1120 to prevent strict-MSS endpoints from stalling.
# NOTE: this 1120 integer is rewritten at boot by toggle.sh from GATEWAY_MSS in .env
# (Gluetun's parser can't interpolate variables)  change GATEWAY_MSS in .env, not this line.
iptables -t mangle -A POSTROUTING -p tcp --tcp-flags SYN,RST SYN -j TCPMSS --set-mss 1120

iptables -t nat -A PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT --to-port 5335
iptables -t nat -A PREROUTING -i tailscale0 -p tcp --dport 53 -j REDIRECT --to-port 5335

# Also redirect DNS from Docker bridge interface so the host can reach AdGuard via localhost:53
iptables -t nat -A PREROUTING -i eth0 -p udp --dport 53 -j REDIRECT --to-port 5335
iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 53 -j REDIRECT --to-port 5335

# Block IPv6 exit-node traffic to prevent leakage (Gluetun/WARP is IPv4-only)
ip6tables -A FORWARD -i tailscale0 -j DROP

# Block DNS-over-TLS (Port 853) to force Android Private DNS to fallback to Port 53
iptables -A FORWARD -i tailscale0 -p tcp --dport 853 -j REJECT --reject-with tcp-reset
iptables -A FORWARD -i tailscale0 -p udp --dport 853 -j REJECT

# Block QUIC (UDP 443) to force Facebook/Google to fallback to TCP.
# UDP bypasses TCP MSS clamping, causing fragmentation. Over weak Wi-Fi (far from gateway),
```

```
# dropping a single fragment destroys the entire packet, stalling image downloads.
iptables -A FORWARD -i tailscale0 -p udp --dport 443 -j REJECT
```

27 scripts/Dockerfile.routing-fix

```
FROM golang:1.22-alpine AS builder
WORKDIR /app
COPY logger/ ./logger/
RUN cd logger && go build -o logger main.go

FROM alpine:3.20
RUN apk add --no-cache iptables iproute2 curl ipset
COPY --from=builder /app/logger/logger /usr/local/bin/nullexit-logger
COPY routing-fix.sh /routing-fix.sh
CMD ["sh", "/routing-fix.sh"]
```

28 scripts/Dockerfile.rule-compiler

```
FROM golang:1.22-alpine AS builder
WORKDIR /app
COPY rule-compiler/ ./rule-compiler/
RUN cd rule-compiler && go build -o sync-rules main.go

FROM alpine:3.20
COPY --from=builder /app/rule-compiler/sync-rules /usr/local/bin/sync-rules
WORKDIR /app
CMD ["sync-rules"]
```

29 scripts/Toggle-Gateway.bat

```
<# :
@echo off
powershell -NoProfile -ExecutionPolicy Bypass -Command "Invoke-Expression ([System.IO.File]::ReadAllText('%~f0
    ') -replace '^(?s).*?#\s*>\s*\r?\n','')"
exit /b
#>
# Toggle-Gateway.bat (Hybrid Batch-PowerShell Toggler)  nullexit Gateway Toggler for Windows
# Mirrors the sophisticated error handling, checks, and automation of toggle.sh

#  1. Administrator Check & Elevation
$isAdmin = ([Security.Principal.WindowsPrincipal][Security.Principal.WindowsIdentity]::GetCurrent()).IsInRole([
    Security.Principal.WindowsBuiltInRole]::Administrator)
if (-not $isAdmin) {
    Write-Host "Requesting Administrator privileges..." -ForegroundColor Yellow
    Start-Process PowerShell -ArgumentList "-NoProfile -ExecutionPolicy Bypass -Command `"Invoke-Expression ([
    System.IO.File]::ReadAllText('$PSCommandPath') -replace '^(?s).*?#\s*>\s*\r?\n','')`"" -Verb RunAs
    Exit
}

# Set console output encoding to UTF8 for clean symbols
[Console]::OutputEncoding = [System.Text.Encoding]::UTF8
Clear-Host

Write-Host "" -ForegroundColor Green
Write-Host "           nullexit Windows Toggler           " -ForegroundColor Green
Write-Host "" -ForegroundColor Green
Write-Host ""

#  2. Helpers
function Reset-HostDNS {
    Write-Host "Restoring default DNS (DHCP/DHCP-assigned DNS)..." -ForegroundColor Cyan
    $activeAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
    foreach ($adapter in $activeAdapters) {
        Write-Host "  -> Resetting DNS on adapter: $($adapter.Name)" -ForegroundColor Gray
        Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ResetServerAddresses -ErrorAction
    SilentlyContinue
    }
}
```

```
35
36 function Hijack-HostDNS ($ip) {
37     Write-Host "Hijacking DNS to route through AdGuard ($ip)..." -ForegroundColor Yellow
38     $activeAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
39     foreach ($adapter in $activeAdapters) {
40         Write-Host "  -> Pointing DNS to $ip on adapter: $($adapter.Name)" -ForegroundColor Gray
41         Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ServerAddresses $ip -ErrorAction
       SilentlyContinue
42     }
43 }
44
45 function Get-HostTailscalePath {
46     $paths = @(
47         "$env:ProgramFiles\Tailscale\tailscale.exe",
48         "$env:ProgramFiles (x86)\Tailscale\tailscale.exe",
49         "tailscale.exe"
50     )
51     foreach ($path in $paths) {
52         if (Test-Path $path) { return $path }
53     }
54     return $null
55 }
56
57 #  3. Detect State
58 # Prevent deadlocks by resetting DNS to default temporarily during checks
59 Reset-HostDNS
60
61 Write-Host "Checking Docker daemon status..." -ForegroundColor Cyan
62 & docker info >$null 2>&1
63 if ($LASTEXITCODE -ne 0) {
64     Write-Host "Docker is not running. Starting Docker Desktop..." -ForegroundColor Yellow
65     $dockerPath = "C:\Program Files\Docker\Docker\Docker Desktop.exe"
66     if (Test-Path $dockerPath) {
67         Start-Process $dockerPath
68         Write-Host "Waiting for Docker daemon to become responsive..." -ForegroundColor Gray
69         $timeout = 60
70         while ($timeout -gt 0) {
71             Start-Sleep -Seconds 2
72             & docker info >$null 2>&1
73             if ($LASTEXITCODE -eq 0) {
74                 Write-Host "Docker is ready!" -ForegroundColor Green
75                 break
76             }
77             $timeout -= 2
78         }
79         if ($timeout -le 0) {
80             Write-Error "Docker failed to start in time. Aborting."
81             Exit 1
82         }
83     } else {
84         Write-Error "Docker Desktop executable not found at standard path. Please start Docker manually."
85         Exit 1
86     }
87 }
88
89 # Detect if gateway is already running
90 $runningWarp = docker compose ps --status running --format json | ConvertFrom-Json -ErrorAction
       SilentlyContinue | Where-Object { $_.Service -eq "warp" }
91 $isGatewayActive = ($null -ne $runningWarp)
92
93 #  4. Execution
94 if ($isGatewayActive) {
95     #  STOP PATH
96     Write-Host "Gateway is currently RUNNING. Disabling now..." -ForegroundColor Yellow
97     Write-Host ""
98
99     # Disconnect host Tailscale from exit node if installed
100     $tsCli = Get-HostTailscalePath
101     if ($null -ne $tsCli) {
102         Write-Host "Disconnecting host Tailscale from exit node..." -ForegroundColor Cyan
103         & $tsCli up --exit-node= --accept-dns=true >$null 2>&1
104     }
105
106     Write-Host "Stopping Docker containers..." -ForegroundColor Cyan
107     docker compose down -t 5
108
109     Reset-HostDNS
110     Write-Host "Gateway has been successfully STOPPED." -ForegroundColor Green
111 } else {
112     #  START PATH
113     Write-Host "Gateway is currently STOPPED. Enabling now..." -ForegroundColor Yellow
```

```
    Write-Host ""

    # Clean up corrupted AdGuardHome configurations
    $confFile = "adguard\conf\AdGuardHome.yaml"
    if (Test-Path $confFile) {
        $fileSize = (Get-Item $confFile).Length
        if ($fileSize -eq 0) {
            Remove-Item $confFile -Force
            Write-Host "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop." -ForegroundColor Yellow
        }
    }

    Write-Host "Starting Docker containers..." -ForegroundColor Cyan
    docker compose up -d

    Write-Host "Waiting for container's Tailscale connection to offer Exit Node..." -ForegroundColor Cyan
    Write-Host "  (This can take up to 60 seconds...)" -ForegroundColor Gray

    $elapsed = 0
    $maxWait = 60
    $resolvedIP = $null

    while ($elapsed -lt $maxWait) {
        $statusJson = docker compose exec -T tailscale tailscale status --json 2>$null
        if ($null -ne $statusJson) {
            $status = $statusJson | ConvertFrom-Json -ErrorAction SilentlyContinue
            if ($null -ne $status -and $null -ne $status.Self) {
                # Check if it has a valid Tailscale IP and is running
                if ($status.Self.Online -and $status.Self.TailscaleIPs.Count -gt 0) {
                    $resolvedIP = $status.Self.TailscaleIPs[0]
                    break
                }
            }
        }
        Start-Sleep -Seconds 2
        $elapsed += 2
    }

    if ($null -ne $resolvedIP) {
        Write-Host "Tailscale IP resolved: $resolvedIP" -ForegroundColor Green
        Write-Host ""

        # Configure host Tailscale client to route through this exit node
        $tsCli = Get-HostTailscalePath
        if ($null -ne $tsCli) {
            Write-Host "Configuring host Tailscale to use Exit Node ($resolvedIP)..." -ForegroundColor Cyan
            & $tsCli up --exit-node=$resolvedIP --accept-dns=false >$null 2>&1
        }

        # Hijack DNS
        Hijack-HostDNS $resolvedIP
        Write-Host ""
        Write-Host "Gateway has been successfully STARTED." -ForegroundColor Green
    } else {
        Write-Host "Warning: Could not resolve gateway Tailscale IP. DNS hijacking skipped." -ForegroundColor Red
        # Fallback to .gateway_ip if it exists
        if (Test-Path ".gateway_ip") {
            $fallbackIP = Get-Content ".gateway_ip" -TotalCount 1
            if (-not [string]::IsNullOrWhiteSpace($fallbackIP)) {
                Write-Host "[Fallback] Using static IP from .gateway_ip: $fallbackIP" -ForegroundColor Yellow
                Hijack-HostDNS $fallbackIP
                Write-Host "Gateway STARTED (with static fallback IP)." -ForegroundColor Green
            }
        } else {
            Write-Host "Gateway started in offline/unauthenticated state. Run setup.sh if this is a fresh setup." -ForegroundColor Yellow
        }
    }
}

Write-Host ""
Write-Host "Press any key to exit..."
[void]$Host.UI.RawUI.ReadKey("NoEcho,IncludeKeyDown")
```

30 scripts/check_circuits.py

```python
#!/usr/bin/env python3
import socket
import os
import re

def query_tor(port, command, password=""):
    try:
        s = socket.socket()
        s.settimeout(2)
        s.connect(('127.0.0.1', port))
        s.sendall(f"AUTHENTICATE \"{password}\"\r\n{command}\r\nQUIT\r\n".encode())
        response = b""
        while True:
            chunk = s.recv(4096)
            if not chunk:
                break
            response += chunk
        return response.decode()
    except Exception as e:
        return f"Error connecting to Tor Control Port: {e}"

def get_ports():
    ports = {}
    ports_file = "/tmp/nullexit-ports.env"
    if os.path.exists(ports_file):
        with open(ports_file, 'r') as f:
            for line in f:
                if line.startswith("export "):
                    # Split at first '=' to handle potential multiple '=' in value
                    parts = line.replace("export ", "").strip().split("=", 1)
                    if len(parts) == 2:
                        key, val = parts
                        try:
                            ports[key] = int(val)
                        except ValueError:
                            ports[key] = val
    return ports

def get_password(ports):
    # 1. Try ports.env
    password = ports.get("TOR_PASSWORD")
    if password:
        return password

    # 2. Try .env
    if os.path.exists(".env"):
        try:
            with open(".env", "r") as f:
                for line in f:
                    if line.startswith("ADGUARD_PASSWORD="):
                        return line.split("=", 1)[1].strip().strip('"\'')
                    if line.startswith("TOR_PASSWORD="):
                        return line.split("=", 1)[1].strip().strip('"\'')
        except Exception:
            pass

    # 3. Fallback
    return "nullexit"

def main():
    ports = get_ports()
    control_port = ports.get("TOR_CONTROL_PORT", 9051)
    password = get_password(ports)

    print(f"Connecting to Tor Control Port on 127.0.0.1:{control_port}...")
    status = query_tor(control_port, "GETINFO circuit-status", password)

    if "Error" in status:
        print(status)
        return

    print("\nActive Tor Circuits:")
    circuits = re.findall(r"(\d+)\s+BUILT\s+(.+?)\s+PURPOSE", status)
    if not circuits:
        print("No active built circuits found yet. Tor might still be bootstrapping.")
        return

```

```
    for circ_id, path_str in circuits:
        print(f"\nCircuit #{circ_id}:")
        nodes = path_str.split(",")
        for idx, node in enumerate(nodes):
            fingerprint = node.split("~")[0].replace("$", "")
            nickname = node.split("~")[1] if "~" in node else "unknown"

            # Fetch node IP
            ns_info = query_tor(control_port, f"GETINFO ns/id/{fingerprint}", password)
            ip = "unknown"
            for line in ns_info.split("\n"):
                if line.startswith("r "):
                    parts = line.split(" ")
                    if len(parts) >= 7:
                        ip = parts[6]
                        break

            # Fetch country
            country = "unknown"
            if ip != "unknown":
                country_info = query_tor(control_port, f"GETINFO ip-to-country/{ip}", password)
                for line in country_info.split("\n"):
                    if line.startswith("250-ip-to-country/"):
                        country = line.split("=")[1].strip().upper()
                        break

            role = "Guard" if idx == 0 else ("Exit" if idx == len(nodes) - 1 else "Middle")
            print(f"  [{role}] {nickname} ({ip}) - Country: {country}")

if __name__ == "__main__":
    main()
```

31 scripts/diagnose-host-leak.sh

```
#!/bin/bash
# scripts/diagnose-host-leak.sh  nullexit host-routing diagnostic
#
# Symptom this addresses: tests like https://www.whatismyip.com on the
# *HOST MAC* return the underlying physical ISP/ASN
# local ISP / campus network") even though the nullexit gateway is active and remote
# tailnet clients correctly egress via Cloudflare WARP. The container and
# the gateway itself are healthy  only the host's own routing is leaking.
#
# This script runs 7 checks in sequence, classifies the result into ONE
# of the three known leak scenarios, writes the full report to a
# timestamped file in the project root, and prints a verdict + a
# ready-to-run fix command on stdout. Pass --fix to apply the matching
# remediation in the same invocation and re-verify egress afterwards.
# See devref.md 10.30 for the deep-dive rationale.
#
# Usage:
#   bash scripts/diagnose-host-leak.sh              # diagnose + write report
#   bash scripts/diagnose-host-leak.sh --fix        # diagnose + fix + re-verify
#   bash scripts/diagnose-host-leak.sh --watch      # full baseline then loop checks every 60s
#   bash scripts/diagnose-host-leak.sh --watch 30   # same, with custom interval (seconds)
#   bash scripts/diagnose-host-leak.sh --help       # this message
#
# Multiple runs produce timestamped files (one per run) so historical
# reports can be diffed. The newest file is the most recent.

set -uo pipefail
# NOTE: errexit (`set -e`) is intentionally NOT enabled. Several checks
# below are EXPECTED to fail and that's diagnostic data (e.g. `tailscale
# status` returns non-zero when the daemon is wedged). Every command
# pipes to a helper that records the exit code without killing the script.

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
source "$SCRIPT_DIR/common.sh"
LOG_FILE="$PROJECT_ROOT/output.log"
REPORT_FILE="$PROJECT_ROOT/host-leak-diagnostic-$TIMESTAMP.txt"

#  Argument parsing
DO_FIX=false
DO_WATCH=false
```

```
WATCH_INTERVAL=60
WATCH_LOG="$PROJECT_ROOT/host-leak-watch.log"
case "${1:-}" in
  --fix) DO_FIX=true ;;
  --watch)
    DO_WATCH=true
    # Optional second argument: custom interval in seconds
    if [ -n "${2:-}" ] && [[ "$2" =~ ^[0-9]+$ ]]; then
      WATCH_INTERVAL="$2"
    fi
    ;;
  --help|-h)
    sed -n '2,31p' "$0"
    exit 0
    ;;
  "")
    : # default: diagnose only
    ;;
  *)
    echo "Unknown argument: $1" >&2
    echo "Run with --help for usage." >&2
    exit 2
    ;;
esac

#  Colours (auto-disabled when stdout is not a tty, e.g. piped)
if [ -t 1 ]; then
  RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
  BOLD='\033[1m'; NC='\033[0m'
else
  RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
fi

#  Output helpers  write to BOTH stdout and the report file
exec 3>> "$REPORT_FILE" || {
  printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
  exit 1
}
section() {
  local title="$1"
  printf '\n%b %s %b\n' "$BOLD" "$title" "$NC" >&3
  printf '\n%b %s %b\n' "$BOLD" "$title" "$NC"
}
line()   { printf '  %b\n' "$1" >&3; printf '  %b\n' "$1"; }
ok()     { line "${GREEN}${NC} $1"; }
warn()   { line "${YELLOW}${NC} $1"; }
fail()   { line "${RED}${NC} $1"; }
note()   { line "(no ${1})"; }

#  Watch-mode header (only shown when entering watch loop)
if [ "$DO_WATCH" = "true" ]; then
  echo ""
  echo ""
  echo " n u l l e x i t   w a t c h   m o d e"
  echo ""
  echo " Interval:        ${WATCH_INTERVAL}s"
  echo " Watch log:       $WATCH_LOG"
  echo ""
else
  echo ""
  echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c"
  echo ""
  echo " Timestamp (UTC): $TIMESTAMP"
  echo " Project root:    $PROJECT_ROOT"
  echo " Report file:     $REPORT_FILE"
  echo " Mode:            $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only (pass --
     fix to remediate)")"
  echo ""
  echo "" >&3
  echo "" >&3
  echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c" >&3
  echo "" >&3
  echo " Timestamp (UTC): $TIMESTAMP" >&3
  echo " Mode:            $([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only")" >&3
fi

#  Always add Homebrew to PATH (same as toggle.sh / recover.sh)
export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"

#  Pure-bash timeout (macOS lacks GNU `timeout`)
run_with_timeout() {
```

```
  local timeout_sec="$1"
  shift
  if [ "$#" -eq 0 ]; then return 1; fi
  "$@" &
  local cmd_pid=$!
  (
    sleep "$timeout_sec"
    if kill -0 "$cmd_pid" 2>> "$LOG_FILE"; then
      kill -9 "$cmd_pid" 2>> "$LOG_FILE" || true
    fi
  ) >> "$LOG_FILE" 2>&1 &
  local watcher_pid=$!
  set +e
  wait "$cmd_pid"
  local exit_status=$?
  set -uo pipefail
  kill "$watcher_pid" 2>> "$LOG_FILE" || true
  return "$exit_status"
}

#
# Quick-check functions (used by --watch loop; return simple status strings)
# Each echoes its result so the caller can capture and compare.
#

quick_warp() {
  # Returns: "on", "off", "unknown"
  local out
  out=$(run_with_timeout 8 curl -s --max-time 6 \
        https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null || echo "")
  local warp
  warp=$(printf '%s' "$out" | awk -F'=' '/^warp=/{print $2; exit}')
  printf '%s' "${warp:-unknown}"
}

quick_ipv6() {
  # Returns: the IPv6 address if leaking, or empty string if clean
  local out
  out=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>/dev/null || echo "")
  if [ -n "$out" ] && printf '%s' "$out" | grep -qE '^[0-9a-fA-F:]+$'; then
    printf '%s' "$out"
  fi
}

quick_default_route() {
  # Returns: "utun" if default route is via Tailscale utun*,
  #          "wifi" if via physical Wi-Fi (192.168.x.x),
  #          "other" otherwise
  local route
  route=$(netstat -rn 2>/dev/null | awk '/^default/ {print $0; exit}')
  if [ -z "$route" ]; then
    printf '%s' "none"
  elif printf '%s' "$route" | grep -qE 'utun[0-9]+'; then
    printf '%s' "utun"
  elif printf '%s' "$route" | grep -qE '^(default|0\.0\.0\.0).*\b192\.168\.'; then
    printf '%s' "wifi"
  else
    printf '%s' "other"
  fi
}

#  Watch loop (runs after baseline diagnostic when --watch is passed)
run_watch_loop() {
  local baseline_warp="$1"
  local baseline_default="$2"
  local baseline_ipv6_leak="$3"
  local interval="$4"

  echo ""
  echo ""
  echo " Baseline established. Monitoring every ${interval}s..."
  echo "  warp:     $baseline_warp"
  echo "  default:  $baseline_default"
  echo ""

  #  Safety: warn if interval is very short (avoids hammering APIs)
  if [ "$interval" -lt 30 ]; then
    printf '%b  Short interval (%ss)  rate-limiting risk on external APIs%b\n' "$YELLOW" "$interval" "$NC"
    printf '  Consider 30s or longer for production use.\n'
  fi
  printf '[%s] WATCH START  warp=%s  default=%s  interval=%ss\n' \
```

```
    "$(date -u +%Y-%m-%dT%H:%M:%SZ)" "$baseline_warp" "$baseline_default" "$interval" \
    >> "$WATCH_LOG"

  local iteration=0
  local last_warp="$baseline_warp"
  local last_default="$baseline_default"
  local last_ipv6_ok=$([ "$baseline_ipv6_leak" = "false" ] && echo true || echo false)

  trap 'printf "\n%bWatch stopped by user. Log at: %s%b\n" "$YELLOW" "$WATCH_LOG" "$NC"; exit 0' INT TERM

  while true; do
    iteration=$((iteration + 1))

    local warp_now ipv6_now default_now
    warp_now=$(quick_warp)
    ipv6_now=$(quick_ipv6)
    default_now=$(quick_default_route)

    local ts
    ts=$(date -u +%H:%M:%S)

    #  WARP check
    if [ "$warp_now" != "$last_warp" ]; then
      if [ "$warp_now" = "off" ] || [ "$warp_now" = "unknown" ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bWARP FLIPPED: %s  %s%b  YOUR IP IS LEAKING!\n' \
          "$RED" "$ts" "$BOLD" "$last_warp" "$warp_now" "$NC"
        printf '[%s] ALERT  warp flipped  %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
      else
        printf '\n%b[%s]  warp recovered: %s  %s%b\n' \
          "$GREEN" "$ts" "$last_warp" "$warp_now" "$NC"
        printf '[%s] OK     warp recovered  %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
      fi
      last_warp="$warp_now"
    fi

    #  IPv6 check
    if [ -n "$ipv6_now" ]; then
      if [ "$last_ipv6_ok" = true ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bIPv6 LEAK DETECTED  %s%b\n' \
          "$RED" "$ts" "$BOLD" "$ipv6_now" "$NC"
        printf '[%s] ALERT  IPv6 leak  %s\n' "$ts" "$ipv6_now" >> "$WATCH_LOG"
        last_ipv6_ok=false
      fi
    else
      if [ "$last_ipv6_ok" = false ]; then
        printf '%b[%s]  IPv6 leak resolved%b\n' "$GREEN" "$ts" "$NC"
        printf '[%s] OK     IPv6 resolved\n' "$ts" >> "$WATCH_LOG"
      fi
      last_ipv6_ok=true
    fi

    #  Default route check
    if [ "$default_now" != "$last_default" ]; then
      if [ "$default_now" != "utun" ]; then
        printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"
        printf '%b[%s] %bDEFAULT ROUTE CHANGED: %s  %s%b  traffic may bypass WARP!\n' \
          "$RED" "$ts" "$BOLD" "$last_default" "$default_now" "$NC"
        printf '[%s] ALERT  route changed  %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
      else
        printf '%b[%s]  default route recovered: %s  %s%b\n' \
          "$GREEN" "$ts" "$last_default" "$default_now" "$NC"
        printf '[%s] OK     route recovered  %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
      fi
      last_default="$default_now"
    fi

    #  Heartbeat (every 10th iteration or if all OK)
    if [ $((iteration % 10)) -eq 0 ]; then
      printf '[%s] #%d  warp=%s  ipv6=%s  route=%s\n' \
        "$ts" "$iteration" "$warp_now" "${ipv6_now:-none}" "$default_now" >> "$WATCH_LOG"
      printf '\r%b[%s] #%d  warp=%s  route=%s  (Ctrl+C to stop)%b' \
        "$GREEN" "$ts" "$iteration" "$warp_now" "$default_now" "$NC"
    fi

    sleep "$interval"
  done
}

#
```

```
# Begin main diagnostic (skipped in non-watch modes if we're just watching)
#

ACTIVE_SERVICE=$(get_active_service)
ENO_SERVICE=$(networksetup -listnetworkserviceorder 2>> "$LOG_FILE" \
              | grep -B 1 "Device: en0" | head -1 \
              | sed -E 's/^\([0-9\*]+\) //' || true)
[ -z "$ENO_SERVICE" ] && ENO_SERVICE="Wi-Fi"

# Resolve the gateway's Tailscale IP.
GATEWAY_TS_IP=""
if [ -f "$PROJECT_ROOT/.gateway_ip" ]; then
  GATEWAY_TS_IP=$(cat "$PROJECT_ROOT/.gateway_ip" 2>> "$LOG_FILE" \
                  | tr -d '\r' | awk 'NR==1{print $1; exit}')
fi
if [ -z "$GATEWAY_TS_IP" ] && command -v docker >> "$LOG_FILE" 2>&1 \
   && docker compose ps --status running 2>/dev/null | grep -q 'tailscale'; then
  GATEWAY_TS_IP=$(run_with_timeout 10 docker compose exec -T tailscale \
                  tailscale ip -4 2>> "$LOG_FILE" \
                  | tr -d '\r' | awk 'NR==1{print $1; exit}')
fi

#  Scenario classification state
SCENARIO=""
WARP_STATUS=""
SOCKS5_ENABLED=false
TS_ON_MESH=false
DEFAULT_VIA_UTUN=false
IPV6_LEAK=false

#
section "1/8  Host Tailscale mesh status"
#
if ! command -v tailscale >> "$LOG_FILE" 2>&1; then
  fail "tailscale CLI not on PATH  install via 'brew install tailscale' then re-run"
  echo "  (a Tailscale daemon is required for the exit-node path to work)"
else
  set +e
  TS_OUT=$(run_with_timeout 5 tailscale status 2>&1)
  TS_EXIT=$?
  set -uo pipefail
  TS_HEAD=$(printf '%s\n' "$TS_OUT" | head -5 | tr -d '\r')
  if [ "$TS_EXIT" -eq 137 ]; then
    fail "tailscaled daemon is wedged/unresponsive (5s timeout)"
    echo "  Fix path: 'sudo brew services restart tailscale'"
  elif printf '%s' "$TS_OUT" | grep -qE 'Logged out|not logged in|expired'; then
    fail "tailscale is logged out / auth expired on this host"
    echo "  Fix path: 'sudo tailscale up' (browser auth once)"
  elif printf '%s' "$TS_OUT" | grep -q '100\.'; then
    ok "Host is on tailnet (100.x.x.x visible in status)"
    TS_ON_MESH=true
    if printf '%s' "$TS_OUT" | grep -q "$GATEWAY_TS_IP"; then
      ok "Gateway $GATEWAY_TS_IP visible in host's peer list"
    elif [ -n "$GATEWAY_TS_IP" ]; then
      warn "Gateway $GATEWAY_TS_IP NOT in host's peer list  possible auth/key mismatch"
    fi
    line " first 5 lines of 'tailscale status' "
    line "$(printf '%s' "$TS_HEAD" | awk '{print "    "$0}')"
  else
    fail "tailscale output unrecognized:"
    line "$(printf '%s' "$TS_OUT" | head -3 | awk '{print "    "$0}')"
  fi
fi

#
section "2/8  Active SOCKS5 proxy on Wi-Fi"
#
SOCKS_OUT=$(networksetup -getsocksfirewallproxy "$ACTIVE_SERVICE" 2>> "$LOG_FILE" \
            || networksetup -getsocksfirewallproxy Wi-Fi 2>> "$LOG_FILE" \
            || echo "Could not query SOCKS5 proxy state")
if printf '%s' "$SOCKS_OUT" | grep -q "Enabled: Yes"; then
  fail "SOCKS5 proxy IS enabled on $ACTIVE_SERVICE  toggle.sh fell back to SOCKS5 mode last run"
  SOCKS5_ENABLED=true
  line "  Full state:"
  line "$(printf '%s' "$SOCKS_OUT" | awk '{print "    "$0}')"
  line "  Browse on macOS ignores system SOCKS5 by default  traffic bypasses WARP."
else
  ok "SOCKS5 proxy is OFF on $ACTIVE_SERVICE  exit-node path should be active"
fi

#
```

```
section "3/8  Egress IP + WARP tunnel verification (definitive test)"
#
CDN_OUT=$(run_with_timeout 8 curl -s --max-time 6 \
          https://www.cloudflare.com/cdn-cgi/trace 2>&1 || echo "(curl failed)")
CDN_WARP=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^warp=/{print $2; exit}')
CDN_IP=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^ip=/{print $2; exit}')
CDN_COLO=$(printf '%s' "$CDN_OUT" | awk -F'=' '/^colo=/{print $2; exit}')
WARP_STATUS="$CDN_WARP"
if [ "$CDN_WARP" = "on" ]; then
  ok "warp=on  tunnel is hot, host traffic IS going through Cloudflare WARP"
  line "  Public IP:   $CDN_IP"
  line "  Cloudflare:   ${CDN_COLO:-unknown colo}"
  line "   whatismyip.com's 'local ISP' result is its ISP-database error, not yours."
elif [ "$CDN_WARP" = "off" ]; then
  fail "warp=off  host traffic is NOT going through WARP"
  line "  Public IP:   $CDN_IP  (NOT Cloudflare)"
  warn "This is the actual leak."
else
  warn "could not determine warp status (curl failed or returned unexpected shape)"
  line "  Raw output: $(printf '%s' "$CDN_OUT" | head -3 | tr '\n' '|')"
fi

#
section "4/8  IPv6 leak probe (bypasses IPv4 exit-node on campus networks)"
#
IPV6_OUT=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
if [ -n "$IPV6_OUT" ] && printf '%s' "$IPV6_OUT" | grep -qE '^[0-9a-fA-F:]+$'; then
  fail "host has live IPv6 egress  $IPV6_OUT (A6 record bypasses WARP)"
  IPV6_LEAK=true
  line "  Tailscale exit-node advertisement is IPv4-only, so IPv6 traffic skips it."
  line "  Quick-fix: 'sudo networksetup -setv6off $ACTIVE_SERVICE'"
else
  ok "no IPv6 egress detected (good  IPv6 won't bypass the tunnel)"
fi

#
section "5/8  Host egress route for real traffic (route -n get 1.1.1.1)"
#
# WHY NOT 'netstat -rn | grep default':
# On macOS, Tailscale does NOT replace the physical default route installed by
# DHCP. Instead it adds a second interface-scoped route bound to the utun
# interface that the kernel prefers for real traffic forwarding. The literal
# "default" entry (192.168.x.x via en0) is left untouched in the table as a
# BSD routing quirk  it's not lying, it's just answering a different question.
# The correct question is: "where does a real destination actually resolve?"
# `route -n get 1.1.1.1` asks the kernel's forwarding logic directly.
ROUTE_GET=$(route -n get 1.1.1.1 2>/dev/null)
ROUTE_IFACE=$(echo "$ROUTE_GET" | awk '/interface:/{print $2}')
ROUTE_GATEWAY=$(echo "$ROUTE_GET" | awk '/gateway:/{print $2}')
if [ -z "$ROUTE_IFACE" ]; then
  warn "route -n get 1.1.1.1 returned no interface  routing table may be empty"
else
  line "  interface: $ROUTE_IFACE  gateway: ${ROUTE_GATEWAY:-n/a}"
  if echo "$ROUTE_IFACE" | grep -qE '^utun[0-9]+'; then
    ok "1.1.1.1 resolves via $ROUTE_IFACE  Tailscale interface-scoped route is winning"
    DEFAULT_VIA_UTUN=true
  elif echo "$ROUTE_IFACE" | grep -qE '^en[0-9]+'; then
    fail "1.1.1.1 resolves via $ROUTE_IFACE (physical Wi-Fi)  exit-node interface route not active"
  else
    warn "1.1.1.1 resolves via unexpected interface: $ROUTE_IFACE"
  fi
fi

#
section "6/8  Last toggle.sh START-relevant lines in output.log (if any)"
#
if [ ! -f "$LOG_FILE" ]; then
  note "output.log"
  line "  toggle.sh was never run from this directory. Run './toggle.sh' first,"
  line "  then re-run this diagnostic."
else
  HITS=$(grep -E 'Exit node enabled|SKIP_EXIT_NODE|Pre-flight|Falling back to SOCKS|Starting|Successfully' \
         "$LOG_FILE" 2>/dev/null | tail -8 || true)
  if [ -z "$HITS" ]; then
    note "matching lines in output.log"
  else
    line "$(printf '%s\n' "$HITS" | awk '{print "    "$0}')"
  fi
fi

#
```

```
section "7/8  Gateway IP we should be pointing at"
#
if [ -n "$GATEWAY_TS_IP" ]; then
  ok "gateway Tailscale IP: $GATEWAY_TS_IP"
else
  warn "could not resolve gateway Tailscale IP"
  line "  (looked in .gateway_ip and tried docker compose exec tailscale ip -4)"
fi

#
section "8/8  Tailscale System Extension (App Store conflict check)"
#
SE_PENDING_UNINSTALL=false
SE_ACTIVE=false
if command -v systemextensionsctl >/dev/null 2>&1; then
  SE_OUT=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale')
  if [ -n "$SE_OUT" ]; then
    line "  Tailscale System Extension(s) detected:"
    line "$(printf '%s' "$SE_OUT" | awk '{print "    "$0}')"
    if printf '%s' "$SE_OUT" | grep -q 'waiting to uninstall'; then
      SE_PENDING_UNINSTALL=true
      warn "System Extension is pending uninstall on next reboot."
    else
      SE_ACTIVE=true
      warn "Active Tailscale System Extension detected. This will conflict with Homebrew Tailscale."
    fi
  else
    ok "no Tailscale System Extension detected"
  fi
else
  note "systemextensionsctl not available (not on macOS or permission restricted)"
fi

#
section "CLASSIFICATION    applies ONE of the four known scenarios"
#

if [ "$WARP_STATUS" = "on" ] && [ "$IPV6_LEAK" = "false" ]; then
  SCENARIO="OK"
elif [ "$SE_PENDING_UNINSTALL" = "true" ]; then
  SCENARIO="SE_PENDING"
elif [ "$SOCKS5_ENABLED" = "true" ]; then
  SCENARIO="A"
elif [ "$IPV6_LEAK" = "true" ]; then
  SCENARIO="B"
elif [ "$DEFAULT_VIA_UTUN" = "false" ] && [ "$TS_ON_MESH" = "true" ]; then
  SCENARIO="C"
else
  SCENARIO="UNKNOWN"
fi

case "$SCENARIO" in
  SE_PENDING)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario SE_PENDING  Tailscale System Extension Pending Uninstall${NC}"
    echo "  A prior App Store Tailscale System Extension is pending uninstall."
    echo "  Since System Integrity Protection (SIP) is enabled, macOS blocks manual"
    echo "  uninstallation of terminated system extensions until the next reboot."
    echo "  Brew-managed tailscaled cannot engage the Network Extension slot until this is resolved."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo shutdown -r now"
    echo "    # After reboot, run:  ./toggle.sh"
    ;;
  A)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario A  SOCKS5 fallback lane${NC}"
    echo "  toggle.sh's pre-flight checks failed last run. The script activated"
    echo "  SOCKS5 + local DNS proxy as a fallback. DNS gets hijacked but"
    echo "  browsers on macOS ignore the system SOCKS5  traffic egresses"
    echo "  direct through the local network."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    echo 'GATEWAY_BYPASS_PING=true' >> $PROJECT_ROOT/.env"
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
    if [ -n "$GATEWAY_TS_IP" ]; then
    echo "        --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
    else
    echo "        --exit-node=\$GATEWAY_TS_IP --exit-node-allow-lan-access=true  # fill in $GATEWAY_TS_IP"
    fi
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
```

```
    echo "    ./toggle.sh"
    ;;
  B)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario B  IPv6 leak over dual-stack campus/local IPv6${NC}"
    echo "  Tailscale exit-node advertises only IPv4. macOS sends AAAA queries"
    echo "  straight out en0, which is dual-stack on most campus APs  IPv6"
    echo "  traffic bypasses the WARP tunnel entirely."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo networksetup -setv6off $ACTIVE_SERVICE"
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
    echo "    # Re-enable IPv6 later with:  sudo networksetup -setv6automatic $ACTIVE_SERVICE"
    ;;
  C)
    echo ""
    echo -e "  ${RED}${BOLD}VERDICT:  Scenario C  Tailscale route-freeze${NC}"
    echo "  Host's default route points at the Wi-Fi gateway instead of the"
    echo "  Tailscale utun* interface. The exit-node preference is set but the"
    echo "  routing-table assertion did not take effect (devref 10.26)."
    echo ""
    echo "  ${BOLD}Recommended fix:${NC}"
    echo "    sudo brew services restart tailscale"
    echo "    sudo route delete -net 0.0.0.0/1"
    echo "    sudo route delete -net 128.0.0.0/1"
    echo "    sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
    if [ -n "$GATEWAY_TS_IP" ]; then
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
    echo "        --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
    else
    echo "    sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true --exit-node=\$TS_IP
     --exit-node-allow-lan-access=true"
    fi
    echo "    ./toggle.sh"
    ;;
  OK)
    echo ""
    echo -e "  ${GREEN}${BOLD}VERDICT:  No host leak detected  tunnel is working${NC}"
    echo "  The local ISP string from whatismyip.com is its IP-database"
    echo "  misclassifying the Cloudflare egress IP.  CDN-CGI's trace endpoint"
    echo "  (which is what we use here) reports warp=on because that endpoint"
    echo "  is hosted by Cloudflare itself and can see its own edge."
    echo "  If you want a third-party confirmation:"
    echo "    curl -s https://api.ipify.org; echo"
    ;;
  *)
    echo ""
    echo -e "  ${YELLOW}${BOLD}VERDICT:  Could not classify automatically${NC}"
    echo "  The diagnostic data did not match any known scenario cleanly."
    echo "  Likely causes:"
    echo "    - tailscaled is logged out / expired (see Section 1)"
    echo "    - Docker is not running (gateway containers down)"
    echo "    - .gateway_ip missing AND docker compose exec failed"
    echo "  Inspect the report at: $REPORT_FILE"
    ;;
esac
echo ""

#
# Write the verdict block to the report file
#
SECOND_BLOCK=$(cat <<EOF


 C L A S S I F I C A T I O N   R E S U L T

 Scenario: $SCENARIO
 WARP egress:  $WARP_STATUS  (cdn-cgi/trace)
 SOCKS5 lane:  $([ "$SOCKS5_ENABLED" = true ] && echo "ENABLED (browsers bypass)" || echo "off")
 IPv6 leak:    $([ "$IPV6_LEAK" = true ] && echo "YES (campus IPv6 egress)" || echo "no")
 Default via:  $([ "$DEFAULT_VIA_UTUN" = true ] && echo "utun* (Tailscale)" || echo "Wi-Fi gateway")
 Host on mesh: $([ "$TS_ON_MESH" = true ] && echo "yes" || echo "no")
 SE pending:   $([ "$SE_PENDING_UNINSTALL" = true ] && echo "yes" || echo "no")
 Gateway IP:   ${GATEWAY_TS_IP:-unresolved}
EOF
)
printf '%s\n' "$SECOND_BLOCK" >&3

#
# --fix mode
#
```

```
if [ "$DO_FIX" = "true" ] && [ "$SCENARIO" != "OK" ] && [ "$SCENARIO" != "UNKNOWN" ]; then
  echo ""
  echo ""
  echo " A P P L Y I N G   F I X  (scenario $SCENARIO)"
  echo ""
  echo ""

  case "$SCENARIO" in
    SE_PENDING)
      warn "System Extension uninstall is pending on next reboot."
      line "Please run 'sudo shutdown -r now' to reboot the system."
      ;;
    A)
      if [ -n "$GATEWAY_TS_IP" ]; then
        line " writing GATEWAY_BYPASS_PING=true to .env"
        ENV="$PROJECT_ROOT/.env"
        [ -f "$ENV" ] || touch "$ENV"
        if grep -q '^GATEWAY_BYPASS_PING=' "$ENV" 2>/dev/null; then
          sed -i '' 's/^GATEWAY_BYPASS_PING=.*/GATEWAY_BYPASS_PING=true/' "$ENV"
        else
          printf '\nGATEWAY_BYPASS_PING=true\n' >> "$ENV"
        fi
        line " re-asserting tailscale exit-node (with --reset)"
        sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
             --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
             >> "$LOG_FILE" 2>&1 \
          && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
          || warn "tailscale up did not return success (check $LOG_FILE)"
      else
        warn "no gateway Tailscale IP resolved  skipping tailscale up"
        warn "fix manually: 'sudo tailscale up --reset ... --exit-node=<TS_IP>'"
      fi
      ;;
    B)
      line " disabling IPv6 on $ACTIVE_SERVICE while gateway is up"
      sudo -n networksetup -setv6off "$ACTIVE_SERVICE" >> "$LOG_FILE" 2>&1 \
        && ok "IPv6 disabled on $ACTIVE_SERVICE (re-enable with 'setv6automatic')" \
        || warn "networksetup -setv6off failed (check $LOG_FILE)"
      ;;
    C)
      line " restarting tailscaled (route-table fix)"
      run_with_timeout 15 brew services restart tailscale >> "$LOG_FILE" 2>&1 \
        || run_with_timeout 15 sudo -n brew services restart tailscale >> "$LOG_FILE" 2>&1
      sleep 3
      line " flushing stale host routes + DNS cache"
      sudo -n route delete -net 0.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
      sudo -n route delete -net 128.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
      sudo -n dscacheutil -flushcache >> "$LOG_FILE" 2>&1 || true
      sudo -n killall -HUP mDNSResponder >> "$LOG_FILE" 2>&1 || true
      if [ -n "$GATEWAY_TS_IP" ]; then
        line " re-asserting exit-node after restart"
        sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
             --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
             >> "$LOG_FILE" 2>&1 \
          && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
          || warn "tailscale up did not return success (check $LOG_FILE)"
      fi
      ;;
  esac

  echo ""
  echo "Re-verifying after fix..."
  sleep 3
  CDN_WARP_AFTER=$(run_with_timeout 8 curl -s --max-time 6 \
      https://www.cloudflare.com/cdn-cgi/trace 2>&1 \
      | awk -F'=' '/^warp=/{print $2; exit}')
  IPV6_AFTER=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
  printf '\n[AFTER FIX]\n' >&3
  printf '  warp status:    %s\n' "$CDN_WARP_AFTER" >&3
  printf '  IPv6 egress:     %s\n' "${IPV6_AFTER:-none}" >&3
  line "  warp status:    ${CDN_WARP_AFTER:-unknown}"
  line "  IPv6 egress:    ${IPV6_AFTER:-none}"
  if [ "$CDN_WARP_AFTER" = "on" ] && [ -z "$IPV6_AFTER" ]; then
    echo ""
    ok "Fix verified  host is now routing through WARP. Remote clients unaffected."
  else
    echo ""
    warn "Fix did not fully take effect on first try. Likely causes:"
    line "   tailscaled needed a longer settling window (re-run diagnostic in 30s)"
    line "   The classified scenario was wrong; paste this report for triage."
    line "   Tailscale needs a full auth refresh: 'sudo tailscale up'"
```

```
  fi
fi

echo ""
echo ""
echo "  Full report written to:"
echo "    $REPORT_FILE"
echo ""
echo ""

#
# --watch mode: enter the continuous monitoring loop after baseline diagnostic
#
if [ "$DO_WATCH" = "true" ]; then
  # Build baseline from the full diagnostic just completed
  BASELINE_WARP="$WARP_STATUS"
  BASELINE_DEFAULT=$([ "$DEFAULT_VIA_UTUN" = "true" ] && echo "utun" || echo "wifi/other")

  if [ "$SCENARIO" != "OK" ]; then
    warn "Baseline diagnostic found issues (scenario: $SCENARIO)."
    line "  Watch loop will still run  alerts fire on any STATE CHANGE from current baseline."
    line "  Fix the issue first? Re-run with:  bash scripts/diagnose-host-leak.sh --fix"
    echo ""
  fi

  run_watch_loop "$BASELINE_WARP" "$BASELINE_DEFAULT" "$IPV6_LEAK" "$WATCH_INTERVAL"
fi

exit 0
```

32 scripts/dns-proxy.py

```
#!/usr/bin/env python3
"""
dns-proxy.py  Local DNS proxy for nullexit gateway.

Listens on UDP port 53, receives DNS queries, forwards them over TCP
to AdGuard at 127.0.0.1:5354 (the Docker port mapping), and returns
the response via UDP. Handles the DNS-over-TCP wire format (2-byte
length prefix) correctly  unlike socat which just forwards raw bytes.
"""

import socket
import struct
import sys
import os
from concurrent.futures import ThreadPoolExecutor

LISTEN_ADDR = os.environ.get("LISTEN_ADDR", "127.0.0.1")
LISTEN_PORT = int(os.environ.get("LISTEN_PORT", 53))
TARGET_HOST = os.environ.get("TARGET_HOST", "127.0.0.1")
TARGET_PORT = int(os.environ.get("TARGET_PORT", 5354))


def handle_query(data: bytes, client_addr: tuple, sock: socket.socket) -> None:
    """Forward a single DNS query over TCP and send the response back via UDP."""
    try:
        tcp = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        tcp.settimeout(5)
        tcp.connect((TARGET_HOST, TARGET_PORT))

        # DNS over TCP: prepend 2-byte length (big-endian)
        tcp.sendall(struct.pack("!H", len(data)) + data)

        # Read the 2-byte response length
        raw_len = tcp.recv(2)
        if len(raw_len) < 2:
            tcp.close()
            return
        resp_len = struct.unpack("!H", raw_len)[0]

        # Read the full response
        response = b""
        while len(response) < resp_len:
            chunk = tcp.recv(resp_len - len(response))
            if not chunk:
```

```
                break
            response += chunk

        tcp.close()

        # Send the response back over UDP (strip TCP length prefix)
        if response:
            sock.sendto(response, client_addr)
    except Exception:
        # Silently drop failed queries (malformed, timeout, etc.)
        pass


def main() -> None:
    # Create UDP socket
    sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
    sock.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)

    try:
        sock.bind((LISTEN_ADDR, LISTEN_PORT))
    except PermissionError:
        print(f"dns-proxy: ERROR  need root to bind port {LISTEN_PORT}", flush=True)
        sys.exit(1)
    except OSError as e:
        print(f"dns-proxy: ERROR  {e}", flush=True)
        sys.exit(1)

    print(f"dns-proxy: listening on UDP {LISTEN_ADDR}:{LISTEN_PORT}  TCP {TARGET_HOST}:{TARGET_PORT}", flush=
     True)

    # Use a thread pool to handle concurrent DNS queries efficiently without thread exhaustion
    executor = ThreadPoolExecutor(max_workers=64)

    while True:
        try:
            data, client_addr = sock.recvfrom(4096)  # Support EDNS0 (up to 4096 bytes) to prevent truncation
            executor.submit(handle_query, data, client_addr, sock)
        except KeyboardInterrupt:
            break
        except Exception:
            pass

    executor.shutdown(wait=False)
    sock.close()


if __name__ == "__main__":
    main()
```

33 scripts/fix-docker-bridge-collision.sh

```
#!/bin/bash
# scripts/fix-docker-bridge-collision.sh  nullexit fix for devref 9.13
#
# Symptom: host default route points at Docker's bridge gateway (172.17.0.1)
# instead of the real Wi-Fi gateway. ALL host internet egress fails. Remote
# tailnet clients route through the gateway correctly (their tunnel is
# independent of the host's broken route table), so they look healthy while
# the host itself cannot reach anything.
#
# Root cause: Docker's default bridge claims 172.17.0.0/16 by default. When the
# host's physical Wi-Fi also lands in the same /16  extremely common on
# university networks which hand out 172.16.0.0/12  the kernel
# installs Docker's bridge as the default route, and every packet bound for
# the internet is silently absorbed by docker0 instead of egressing.
#
# This script:
#   1. Detects the actual collision (Docker's 172.17.0.0/16 vs the host's
#      Wi-Fi subnet, queried live via the ACTIVE interface  not hardcoded en0)
#   2. Picks a non-colliding Docker bip using /8-family routing (works for
#      172.16.0.0/12 campus networks; naive prefix-match would pick 172.26/24
#      which still lives INSIDE a /12 of 172.16.0.0/12)
#   3. Backs up ~/.colima/default/colima.yaml with an ISO-8601 timestamp
#   4. Edits the file in place, idempotent (replace existing docker.bip /
#      append new docker.bip without disturbing other keys; skip YAML comments)
#   5. Restarts Colima so the VM rebuilds with the new bridge subnet
```

```
#   6. Rebinds Wi-Fi DHCP so the host re-learns its real Wi-Fi gateway
#   7. Verifies the new default route is no longer the Docker bridge
#   8. Re-runs toggle.sh to bring the gateway back up cleanly
#   9. Re-runs scripts/diagnose-host-leak.sh and prints the new verdict
#
# Usage:
#   bash scripts/fix-docker-bridge-collision.sh        # apply the fix
#   bash scripts/fix-docker-bridge-collision.sh --dry  # show what would change, then exit
#   bash scripts/fix-docker-bridge-collision.sh --skip-toggle  # fix but don't restart gateway
#   bash scripts/fix-docker-bridge-collision.sh --help # usage

set -uo pipefail
# NOTE: errexit (`set -e`) is intentionally NOT enabled so intermediate
# `command  || warn ` patterns can continue past non-fatal failures.
# Every critical failure (colima restart, sed, cp backup, toggle.sh)
# uses an explicit `|| die` so a half-applied state never escapes.

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
LOG_FILE="$PROJECT_ROOT/output.log"
source "$SCRIPT_DIR/common.sh"
TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
COLIMA_YAML="$HOME/.colima/default/colima.yaml"
PLATFORM="$(uname -s)"

#  Argument parsing
DRY_RUN=false
SKIP_TOGGLE=false
case "${1:-}" in
  --dry|-n)       DRY_RUN=true ;;
  --skip-toggle)  SKIP_TOGGLE=true ;;
  --help|-h)
    sed -n '28,37p' "$0"
    exit 0
    ;;
  "")
    : # default: apply
    ;;
  *)
    printf 'Unknown argument: %s\n' "$1" >&2
    exit 2
    ;;
esac



# Helper: print "[dry-run] $cmd" or actually run $cmd. Wraps the redirect
# INSIDE the function so outer `>> $FILE` statements never accidentally
# append dry-run noise to the real file. (Earlier draft didn't and the
# dry-run polluted $COLIMA_YAML.)
run() {
  if [ "$DRY_RUN" = "true" ]; then
    printf '    [dry-run] %s\n' "$*"
    return 0
  fi
  printf '    $ %s\n' "$*"
  "$@"
}

# Platform-aware sed in-place. macOS BSD sed needs an empty argument;
# GNU sed (Linux) does not. Earlier draft used `-i ''` everywhere and
# would have failed on Linux silently.
sed_inplace() {
  local expr="$1"
  local file="$2"
  case "$PLATFORM" in
    Darwin) run sed -i '' "$expr" "$file" ;;
    *)      run sed -i    "$expr" "$file" ;;
  esac
}

#  0. Resolve the ACTIVE physical interface once (used everywhere)
# Replaces the earlier `en0` hardcoding. Same algorithm recover.sh uses:
# `route get default`  if utun/tun, walk en0..en5 first with `status: active`
# + `inet`  fall back to en0. This means the script also works on hosts
# whose primary NIC is en1/en2 (Thunderbolt Ethernet, USB-LAN, etc.).
resolve_active_iface() {
  local iface
  iface=$(route get default 2>> "$LOG_FILE" | awk '/interface:/{print $2; exit}')
  if [ -z "$iface" ] || [ "${iface#utun}" != "$iface" ] || [ "${iface#tun}" != "$iface" ]; then
    for i in en0 en1 en2 en3 en4 en5; do
```

```
      if ifconfig "$i" 2>> "$LOG_FILE" | grep -q "status: active" \
         && ifconfig "$i" 2>> "$LOG_FILE" | grep -q "inet "; then
        iface="$i"; break
      fi
    done
  fi
  [ -z "$iface" ] && iface="en0"
  printf '%s\n' "$iface"
}
ACTIVE_IFACE=$(resolve_active_iface)

#  1. Detect collision
step "1/9  Detecting Docker-bridge vs Wi-Fi subnet collision"

# Host subnet on the ACTIVE interface (not hardcoded en0).
HOST_SUBNET=$(ifconfig "$ACTIVE_IFACE" 2>> "$LOG_FILE" \
              | awk '/inet /{print $2; exit}')
HOST_GATEWAY=$(route get default 2>> "$LOG_FILE" \
               | awk '/gateway:/{print $2; exit}')
DEFAULT_DEV=$(route get default 2>> "$LOG_FILE" \
              | awk '/interface:/{print $2; exit}')

printf '  active iface:  %s\n' "$ACTIVE_IFACE"
printf '  iface inet:    %s\n' "${HOST_SUBNET:-none}"
printf '  default via:   %s\n' "${HOST_GATEWAY:-none}"
printf '  default dev:   %s\n' "${DEFAULT_DEV:-none}"

# Collision fires when the host's default-route gateway is 172.17.0.1 (Docker's
# bridge) OR when the host subnet is in 172.17.0.0/16. The full-bridge-collision
# case (gateway is literally the docker bridge IP) is the smoking gun.
COLLISION=false
if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#172.17.}" != "$HOST_GATEWAY" ]; then
  COLLISION=true
  fail "default gateway $HOST_GATEWAY collides with Docker's default bridge (172.17.0.0/16)"
elif [ -n "$HOST_SUBNET" ] && [ "${HOST_SUBNET#172.17.}" != "$HOST_SUBNET" ]; then
  COLLISION=true
  fail "host subnet $HOST_SUBNET overlaps Docker's 172.17.0.0/16 bridge"
else
  ok "no collision detected between host Wi-Fi and Docker's default bridge"
fi

if [ "$COLLISION" = "false" ]; then
  echo ""
  echo "Nothing to fix. If you still see egress issues, run:"
  echo "    bash scripts/diagnose-host-leak.sh"
  exit 0
fi

#  2. Pick non-colliding bip via /8-family routing
step "2/9  Picking a non-colliding Docker bip for ~/.colima/default/colima.yaml"

# Earlier draft tried naive prefix-match: "if candidate's 3-octet prefix is
# not a literal prefix of host gateway, pick it". That breaks on /12 campus
# networks (172.16.0.0/12 covers 172.16-172.31): a candidate like
# 172.26.0.1/24 LIVES INSIDE the host's /12 even though the prefix test
# passes. Fix: classify the host's address family FIRST, then choose ALL
# candidates from a DIFFERENT RFC1918 family. /8 boundaries are coarse
# enough that no /24 picked this way can ever collide.
case "$HOST_GATEWAY" in
  172.16.*|172.17.*|172.18.*|172.19.*|172.20.*|172.21.*|172.22.*|172.23.*|\
  172.24.*|172.25.*|172.26.*|172.27.*|172.28.*|172.29.*|172.30.*|172.31.*)
    HOST_FAMILY="172"
    # Host is in 172.x  jump past 172.31 entirely. Pick from 10.x / 192.168.x.
    CANDIDATES="10.200.0.1/24 10.201.0.1/24 192.168.99.1/24"
    ;;
  10.*)
    HOST_FAMILY="10"
    # Host is in 10.x  pick from 172.x (well outside any commonly-deployed
    # 10.0.0.0/8 slice) / 192.168.x.
    CANDIDATES="172.26.0.1/24 172.28.0.1/24 192.168.99.1/24"
    ;;
  192.168.*)
    HOST_FAMILY="192"
    CANDIDATES="172.26.0.1/24 10.200.0.1/24 10.201.0.1/24"
    ;;
  *)
    HOST_FAMILY="unknown"
    warn "host gateway $HOST_GATEWAY is not in any RFC1918 range (/8 family)"
    warn "falling back to abstract candidate set; verify collision-free manually"
    CANDIDATES="172.26.0.1/24 10.200.0.1/24 192.168.99.1/24"
    ;;
```

```
esac

printf '  host address family: %s\n' "$HOST_FAMILY"
printf '  candidate pool:      %s\n' "$CANDIDATES"

# Sanity-check the chosen pool has at least one entry that doesn't literally
# share a /24 with the host gateway. Naive check (just /24 prefix) is fine
# here because we already picked across families.
CHOSEN=""
for cand in $CANDIDATES; do
  PREFIX3=$(printf '%s' "$cand" | cut -d. -f1-3)
  if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#"$PREFIX3".}" != "$HOST_GATEWAY" ]; then
    warn "skip $cand (literal /24 collision with host gateway $HOST_GATEWAY)"
    continue
  fi
  CHOSEN="$cand"
  ok "chose bip=$CHOSEN (cross-family, no /24 literal collision)"
  break
done

if [ -z "$CHOSEN" ]; then
  CHOSEN="172.26.0.1/24"
  warn "could not find a clean candidate; defaulting to $CHOSEN"
  warn "verify with 'route get default' after restart"
fi

#  3. Backup colima.yaml
step "3/9  Backing up $COLIMA_YAML"
mkdir -p "$(dirname "$COLIMA_YAML")"
BACKUP="$HOME/.colima/default/colima.yaml.bak.$TIMESTAMP"
if [ -f "$COLIMA_YAML" ]; then
  # `cp` is safe in dry-run because run() suppresses the inner command
  # entirely  no chance of accidentally writing a real backup.
  run cp -p "$COLIMA_YAML" "$BACKUP" || die "backup failed"
  if [ "$DRY_RUN" = "false" ]; then ok "backup: $BACKUP"; fi
else
  warn "no existing $COLIMA_YAML (will create fresh)"
fi

#  4. Edit colima.yaml in place (idempotent)
step "4/9  Setting docker.bip=$CHOSEN in colima.yaml (in place)"

# Resolve the existing docker.bip (if any) without picking up commented-out
# YAML keys. awk-based extraction; the negated-comment check (`! /^[[:space:]]*#`)
# ensures `# bip: 1.2.3.4/24` is treated as a comment, not the active value.
EXISTING_BIP=""
if [ -f "$COLIMA_YAML" ]; then
  EXISTING_BIP=$(awk '
    /^[[:space:]]*#/ {next}
    /^[[:space:]]*bip:[[:space:]]*/ {print $2; exit}
  ' "$COLIMA_YAML" 2>> "$LOG_FILE" || true)
fi

if [ -n "$EXISTING_BIP" ] && [ "$EXISTING_BIP" = "$CHOSEN" ]; then
  ok "colima.yaml already has docker.bip=$CHOSEN  no edit needed"
elif [ -n "$EXISTING_BIP" ]; then
  printf '  (replacing existing docker.bip=%s with %s)\n' "$EXISTING_BIP" "$CHOSEN"
  # Use sed_inplace so Linux / macOS get the right syntax. Skip YAML comment
  # lines so existing `# bip:` comments aren't rewired.
  sed_inplace "s|^\\([[:space:]]*bip:[[:space:]]*\\).*\$|\\1$CHOSEN|" "$COLIMA_YAML" \
    || die "sed replacement failed (colima.yaml malformed?)"
  ok "colima.yaml updated"
else
  # No existing `bip:` key. Branch:
  #   - file has `docker:` block  append `  bip:` under it
  #   - file has no `docker:` block  append new docker: section
  #   - file does not exist  create fresh
  # Each branch is dry-run-aware so no branch can accidentally leak noise
  # into $COLIMA_YAML during a preview.
  HAS_DOCKER_BLOCK=false
  if [ -f "$COLIMA_YAML" ] && grep -q '^docker:' "$COLIMA_YAML"; then
    HAS_DOCKER_BLOCK=true
  fi

  if [ "$HAS_DOCKER_BLOCK" = "true" ]; then
    if [ "$DRY_RUN" = "true" ]; then
      printf '    [dry-run] would append to %s:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
    else
      printf '\n  bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
        || die "append failed"
      ok "appended bip under existing docker: block"
```

```
    fi
  elif [ -f "$COLIMA_YAML" ]; then
    if [ "$DRY_RUN" = "true" ]; then
      printf '    [dry-run] would append to %s:\\ndocker:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
    else
      printf '\ndocker:\n  bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
        || die "append failed"
      ok "appended new docker: block at end of file"
    fi
  else
    if [ "$DRY_RUN" = "true" ]; then
      printf '    [dry-run] would create %s with:\\ndocker:\\n  bip: %s\\n' "$COLIMA_YAML" "$CHOSEN"
    else
      printf 'docker:\n  bip: %s\n' "$CHOSEN" > "$COLIMA_YAML" \
        || die "create failed"
      ok "created fresh $COLIMA_YAML with docker.bip=$CHOSEN"
    fi
  fi
fi

# Always show the final colima.yaml docker section for visibility
printf '\n%b  colima.yaml `docker:` section now reads:%b\n' "$BOLD" "$NC"
awk '
  /^docker:/{p=1}
  p && /^[^ ]/ && !/^docker:/{exit}
  p
' "$COLIMA_YAML" 2>/dev/null | sed 's/^/    /' || true

if [ "$DRY_RUN" = "true" ]; then
  printf '\n  %b(dry-run: exiting before colima restart / dhcp rebind / toggle.sh)%b\n' "$YELLOW" "$NC"
  exit 0
fi

#  5. Restart Colima
step "5/9  Restarting Colima VM with new bip"
if ! command -v colima >> "$LOG_FILE" 2>&1; then
  die "colima not on PATH  install with 'brew install colima' first"
fi
# No timeout wrapper (would mask actual VM startup time). If colima hangs,
# the user kills with Ctrl+C and can manually `colima restart` from a fresh
# terminal  the rest of this script is idempotent so re-running works.
colima restart >> "$LOG_FILE" 2>&1 || die "colima restart failed; check $LOG_FILE"
ok "colima restarted"

#  6. Rebind DHCP on the ACTIVE interface (not hardcoded en0)
step "6/9  Rebinding DHCP on $ACTIVE_IFACE so host re-learns its real gateway"

# Map ifname  macOS network service name (Wi-Fi, USB 10/100 LAN, etc.).
# Same logic as the diagnostic + toggle.sh.
ACTIVE_SVC=""
if [ "$PLATFORM" = "Darwin" ]; then
  ACTIVE_SVC=$(get_active_service)
fi

case "$PLATFORM" in
  Darwin)
    sudo -n networksetup -setdhcp "$ACTIVE_SVC" renew >> "$LOG_FILE" 2>&1 \
      && ok "DHCP rebind issued on $ACTIVE_SVC" \
      || warn "setdhcp renew failed; try 'sudo ipconfig set $ACTIVE_IFACE DHCP' manually"
    ;;
  Linux)
    sudo -n dhclient -r "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 || true
    sudo -n dhclient "$ACTIVE_IFACE"     >> "$LOG_FILE" 2>&1 \
      && ok "dhclient rebind on $ACTIVE_IFACE" \
      || warn "dhclient rebind failed"
    ;;
  *)
    warn "unsupported platform $PLATFORM; skip DHCP rebind"
    ;;
esac

#  7. Verify new default route
step "7/9  Verifying the new default route is NOT the Docker bridge"
sleep 2  # give DHCP + kernel a beat
NEW_DEFAULT=$(netstat -rn 2>> "$LOG_FILE" | awk '/^default/ {print $0; exit}')
printf '  %s\n' "${NEW_DEFAULT:-no default route detected}"
NEW_GATEWAY=$(printf '%s' "$NEW_DEFAULT" | awk '{print $2; exit}')

if [ -n "$NEW_GATEWAY" ] && [ "${NEW_GATEWAY#172.17.}" != "$NEW_GATEWAY" ]; then
  fail "default route STILL points at 172.17.x.x  DHCP rebind may need manual touch"
  warn "try: sudo ipconfig set $ACTIVE_IFACE DHCP  (or reconnect Wi-Fi in the menu bar)"
```

```
elif [ -n "$NEW_GATEWAY" ]; then
  ok "default route is now via $NEW_GATEWAY (no longer Docker's bridge)"
else
  warn "could not determine new gateway; verify with 'route get default'"
fi

#  8. Re-run toggle.sh to bring the gateway back up
if [ "$SKIP_TOGGLE" = "true" ]; then
  step "8/9  Skipping toggle.sh (--skip-toggle)"
  warn "gateway is currently down  run ./toggle.sh manually when ready"
else
  step "8/9  Re-running toggle.sh to bring the gateway back up"
  if bash "$PROJECT_ROOT/toggle.sh" 2>&1 | tee -a "$LOG_FILE"; then
    ok "toggle.sh completed"
  else
    warn "toggle.sh returned non-zero  see $LOG_FILE for details"
  fi
fi

#  9. Re-run the diagnostic + show verdict
step "9/9  Re-running host-leak diagnostic to confirm fix"
if bash "$SCRIPT_DIR/diagnose-host-leak.sh" 2>&1 | tee -a "$LOG_FILE"; then
  LATEST_REPORT=$(ls -t "$PROJECT_ROOT"/host-leak-diagnostic-*.txt 2>/dev/null | head -1)
  if [ -n "$LATEST_REPORT" ]; then
    VERDICT=$(awk '/^ Scenario:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
    WARP=$(awk -F': *' '/^ WARP egress:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
    printf '\n  Most recent verdict: Scenario=%s  WARP=%s\n' \
           "${VERDICT:-?}" "${WARP:-?}"
    if [ "$VERDICT" = "OK" ]; then
      ok "fix confirmed by diagnostic  host traffic is going through WARP"
    fi
  fi
else
  warn "diagnostic exited non-zero; verify manually"
fi

printf '\n%b%b\n' "$BOLD" "$NC"
printf '%bDone.%b Full transcript at %s\n' "$BOLD" "$NC" "$LOG_FILE"
printf '%b%b\n' "$BOLD" "$NC"
exit 0
```

34 scripts/fix-ssh-delay.sh

```
#!/bin/bash
# scripts/fix-ssh-delay.sh  Fix ~20s SSH connection delay caused by UseDNS
#
# macOS OpenSSH defaults UseDNS to "yes". When Tailscale SSH connects from
# a peer (phone/laptop), sshd performs a reverse-DNS (PTR) lookup on the
# client's 100.x.x.x Tailscale IP. Since DNS is hijacked to the nullexit
# gateway  AdGuard  Cloudflare WARP, and 100.64.0.0/10 is CGNAT private
# space, the PTR query times out (~15-20s) before SSH proceeds.
#
# This script uncomments "UseDNS no" in /etc/ssh/sshd_config and reloads
# sshd. Existing SSH sessions are NOT dropped.
#
# Usage:
#   sudo bash scripts/fix-ssh-delay.sh

set -euo pipefail

SSHD_CONFIG="/etc/ssh/sshd_config"

#  Check we're root
if [ "$(id -u)" -ne 0 ]; then
    echo "ERROR: This script must be run with sudo (it edits $SSHD_CONFIG)"
    echo "Usage: sudo bash scripts/fix-ssh-delay.sh"
    exit 1
fi

#  Apply config (idempotent)
if grep -qE '^UseDNS no' "$SSHD_CONFIG" 2>/dev/null; then
    echo " UseDNS is already set to 'no' in $SSHD_CONFIG"
elif grep -qE '^#UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
    echo " Uncommenting 'UseDNS no' in $SSHD_CONFIG"
    sed -i '' 's/^#UseDNS no/UseDNS no/' "$SSHD_CONFIG"
elif grep -qEi '^UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
```

```
    echo " Changing existing UseDNS line to 'UseDNS no'"
    sed -i '' 's/^UseDNS.*/UseDNS no/' "$SSHD_CONFIG"
else
    echo " Appending 'UseDNS no' to $SSHD_CONFIG"
    echo "" >> "$SSHD_CONFIG"
    echo "UseDNS no" >> "$SSHD_CONFIG"
fi

#  Verify
if ! grep -qE '^UseDNS no' "$SSHD_CONFIG"; then
    echo " Failed to apply UseDNS no  check $SSHD_CONFIG manually"
    exit 1
fi

#  Reload sshd (does NOT drop existing connections)
echo " Reloading sshd..."

SSHD_PID=$(pgrep -f /usr/sbin/sshd 2>/dev/null | head -1 || true)
if [ -n "$SSHD_PID" ]; then
    kill -HUP "$SSHD_PID"
    echo " Sent SIGHUP to sshd (PID $SSHD_PID)  config reloaded"
else
    echo "  No running sshd found. macOS starts sshd on-demand per connection."
    echo "  The new 'UseDNS no' config will take effect on the next SSH connection."
fi

echo ""
echo " Done. SSH connections should now connect instantly."
echo "  Existing sessions were not interrupted."
```

35 scripts/generate_tex.py

```
#!/usr/bin/env python3
import os

PROJECT_ROOT = os.path.abspath(os.path.join(os.path.dirname(__file__), '..'))

IGNORE_DIRS = {'.git', '.github', 'adguard', '__pycache__', 'Recover Gateway.app', 'Toggle Gateway.app'}
IGNORE_FILES = {'.env', 'nullexit_unified.tex', 'nullexit_unified.pdf', 'output.log', 'blocked.log', '.DS_Store
    ', '.lan_p2p_detected', '.signatures', '.gateway_ip', 'TUNNEL_FAILED_CLOSED.marker','.host_ips','refactor.
    md','.build_hash','.rules_hash'}
IGNORE_EXTS = {'.pdf', '.tex', '.png', '.jpg', '.jpeg', '.zip', '.tar', '.gz', '.log', '.aux', '.fls', '.
    fdb_latexmk', '.out', '.toc', '.xdv', '.synctex(busy)'}

# Exclude from code block inclusion, but keep in tree
SKIP_CONTENT_FILES = {'LICENSE'}

PRIORITY_FILES = [
    'README.md',
    'agent.md',
    'devref.md',
    'diagrams.md',
    'toggle.sh',
    'recover.sh',
    'setup.sh',
    'scripts/setup-linux.sh',
    'docker-compose.yml',
    'scripts/common.sh',
    'scripts/watcher.sh',
    'scripts/crypto.sh'
]

def get_language(filename):
    ext = os.path.splitext(filename)[1].lower()
    if ext == '.sh' or ext == '.applescript':
        return 'bash'
    elif ext == '.py':
        return 'Python'
    elif ext == '.go':
        return 'Go'
    elif ext in {'.yml', '.yaml'}:
        return 'bash'
    elif ext == '.plist':
        return 'XML'
    elif ext == '.md':
        return ''
```

```
    return ''

def escape_tex(text):
    return text.replace('_', '\\_')

def generate_tree(dir_path, prefix=""):
    tree_str = ""
    entries = sorted(os.listdir(dir_path))
    entries = [e for e in entries if e not in IGNORE_DIRS and e not in IGNORE_FILES and os.path.splitext(e)[1].
     lower() not in IGNORE_EXTS]

    for i, entry in enumerate(entries):
        path = os.path.join(dir_path, entry)
        is_last = (i == len(entries) - 1)
        connector = "`-- " if is_last else "|-- "
        tree_str += f"{prefix}{connector}{entry}\n"

        if os.path.isdir(path):
            extension = "    " if is_last else "|   "
            tree_str += generate_tree(path, prefix + extension)

    return tree_str

def main():
    tex_path = os.path.join(PROJECT_ROOT, 'nullexit_unified.tex')

    preamble = r"""\documentclass{article}
\usepackage[utf8]{inputenc}
\usepackage[margin=1in]{geometry}
\usepackage{listings}
\usepackage{xcolor}
\usepackage{hyperref}
\usepackage{tocloft}

\definecolor{codegreen}{rgb}{0,0.6,0}
\definecolor{codegray}{rgb}{0.5,0.5,0.5}
\definecolor{codepurple}{rgb}{0.58,0,0.82}
\definecolor{backcolour}{rgb}{0.97,0.97,0.96}

\lstdefinestyle{mystyle}{
    backgroundcolor=\color{backcolour},
    commentstyle=\color{codegreen},
    keywordstyle=\color{blue}\bfseries,
    numberstyle=\tiny\color{codegray},
    stringstyle=\color{codepurple},
    basicstyle=\ttfamily\scriptsize,
    breakatwhitespace=false,
    breaklines=true,
    captionpos=b,
    keepspaces=true,
    numbers=left,
    numbersep=5pt,
    showspaces=false,
    showstringspaces=false,
    showtabs=false,
    tabsize=2,
    frame=single,
    rulecolor=\color{codegray},
    extendedchars=true,
    literate={{-}1 {}{=}1 {}{{OK}}2 {}{-}1 {}{|}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1 {}{+}1
     {}{+}1 {}{v}1
}
\lstset{style=mystyle}

\title{Nullexit: Unified Project Document}
\author{Omar Alaaeldein}
\date{\today}

\begin{document}
\maketitle

\tableofcontents
\newpage

\section{Project Directory Tree}
\begin{verbatim}
nullexit/
"""

    tree = generate_tree(PROJECT_ROOT)
    midamble = r"""\end{verbatim}
```

```
\newpage
"""

    all_relpaths = []
    for root, dirs, files in os.walk(PROJECT_ROOT):
        dirs[:] = sorted([d for d in dirs if d not in IGNORE_DIRS])
        for file in sorted(files):
            if file in IGNORE_FILES or file in SKIP_CONTENT_FILES or os.path.splitext(file)[1].lower() in
     IGNORE_EXTS:
                continue

            filepath = os.path.join(root, file)
            relpath = os.path.relpath(filepath, PROJECT_ROOT)
            all_relpaths.append(relpath)

    def sort_key(path):
        if path in PRIORITY_FILES:
            return (0, PRIORITY_FILES.index(path))
        return (1, path)

    all_relpaths.sort(key=sort_key)

    with open(tex_path, 'w', encoding='utf-8') as f:
        f.write(preamble)
        f.write(tree)
        f.write(midamble)

        for relpath in all_relpaths:
            lang = get_language(relpath)
            lang_attr = f"[language={lang}]" if lang else ""

            f.write(f"\\section{{{escape_tex(relpath)}}}\n")
            f.write(f"\\begin{{lstlisting}}{lang_attr}\n")
            try:
                with open(os.path.join(PROJECT_ROOT, relpath), 'r', encoding='utf-8', errors='ignore') as src:
                    content = src.read()
                    import re
                    # Strip all non-ASCII characters (emojis, box drawings, em dashes, etc)
                    content = re.sub(r'[^\x00-\x7F]+', '', content)
                    f.write(content)
                    if not content.endswith('\n'):
                        f.write('\n')
            except Exception as e:
                f.write(f"Error reading file: {e}\n")
            f.write("\\end{lstlisting}\n\n")

        f.write("\\end{document}\n")

    print(f"Generated {tex_path}")

if __name__ == '__main__':
    main()
```

36 scripts/logger/go.mod

```
module nullexit/logger

go 1.22
```

37 scripts/logger/main.go

```
package main

import (
  "bufio"
  "encoding/json"
  "fmt"
  "io"
  "log"
  "os"
  "regexp"
  "strings"
```

```
  "time"
)

const (
  queryLogPath  = "/adguard_work/data/querylog.json"
  procKmsgPath  = "/proc/kmsg"
  outputLogPath = "/app/blocked.log"
)

func logMessage(msg string) {
  timestamp := time.Now().Format("2006-01-02 15:04:05")
  formattedMsg := fmt.Sprintf("%s %s\n", timestamp, msg)
  fmt.Print(formattedMsg)

  f, err := os.OpenFile(outputLogPath, os.O_APPEND|os.O_CREATE|os.O_WRONLY, 0644)
  if err != nil {
    log.Printf("Error writing to log file: %v\n", err)
    return
  }
  defer f.Close()
  f.WriteString(formattedMsg)
}

func tailFile(filepath string, out chan<- string) {
  for {
    if _, err := os.Stat(filepath); os.IsNotExist(err) {
      time.Sleep(1 * time.Second)
      continue
    }
    break
  }

  f, err := os.Open(filepath)
  if err != nil {
    log.Printf("Error opening file %s: %v\n", filepath, err)
    return
  }
  defer func() {
    if f != nil {
      f.Close()
    }
  }()

  // Seek to end
  f.Seek(0, io.SeekEnd)
  reader := bufio.NewReader(f)

  for {
    line, err := reader.ReadString('\n')
    if err != nil {
      if err == io.EOF {
        stat, err := os.Stat(filepath)
        if err == nil {
          currentOffset, _ := f.Seek(0, io.SeekCurrent)
          if stat.Size() < currentOffset {
            // File truncated or rotated
            f.Close()
            for {
              if _, err := os.Stat(filepath); os.IsNotExist(err) {
                time.Sleep(1 * time.Second)
                continue
              }
              break
            }
            f, _ = os.Open(filepath)
            reader = bufio.NewReader(f)
            continue
          }
        }
        time.Sleep(500 * time.Millisecond)
        continue
      }
      log.Printf("Error reading file %s: %v\n", filepath, err)
      time.Sleep(1 * time.Second)
      continue
    }
    out <- line
  }
}

type adguardLogEntry struct {
```

```
  IP     string `json:"IP"`
  QH     string `json:"QH"`
  QT     string `json:"QT"`
  Result struct {
    IsFiltered bool `json:"IsFiltered"`
    Reason     int  `json:"Reason"`
    Rules      []struct {
      Text string `json:"Text"`
    } `json:"Rules"`
  } `json:"Result"`
}

func monitorDNS() {
  logMessage("[System] Starting DNS block logger...")
  lines := make(chan string)
  go tailFile(queryLogPath, lines)

  reasonMap := map[int]string{
    2: "CustomRule",
    3: "BlockList",
    4: "SafeBrowsing",
    5: "ParentalControl",
    6: "SafeSearch",
    7: "BlockedService",
  }

  for line := range lines {
    var data adguardLogEntry
    if err := json.Unmarshal([]byte(line), &data); err != nil {
      continue
    }

    res := data.Result
    if res.IsFiltered || (res.Reason != 0 && res.Reason != 1) {
      qh := data.QH
      if qh == "" {
        qh = "unknown"
      }
      qt := data.QT
      if qt == "" {
        qt = "unknown"
      }
      clientIP := data.IP
      if clientIP == "" {
        clientIP = "unknown"
      }

      ruleText := "unknown rule"
      if len(res.Rules) > 0 && res.Rules[0].Text != "" {
        ruleText = res.Rules[0].Text
      }

      reasonStr, ok := reasonMap[res.Reason]
      if !ok {
        reasonStr = fmt.Sprintf("Reason-%d", res.Reason)
      }

      logMessage(fmt.Sprintf("[DNS] Blocked %s (Type: %s) for client %s | Reason: %s | Rule: %s", qh, qt,
     clientIP, reasonStr, ruleText))
    }
  }
}

func monitorIPs() {
  logMessage("[System] Starting IP block logger...")

  prefixRe := regexp.MustCompile(`(IP_BLOCK_[A-Z_]+):`)
  srcRe := regexp.MustCompile(`SRC=([0-9a-fA-F\.:]+)`)
  dstRe := regexp.MustCompile(`DST=([0-9a-fA-F\.:]+)`)
  protoRe := regexp.MustCompile(`PROTO=([A-Z0-9]+)`)
  sptRe := regexp.MustCompile(`SPT=(\d+)`)
  dptRe := regexp.MustCompile(`DPT=(\d+)`)
  inRe := regexp.MustCompile(`IN=([a-zA-Z0-9\.\-_]+)`)
  outRe := regexp.MustCompile(`OUT=([a-zA-Z0-9\.\-_]+)`)

  f, err := os.Open(procKmsgPath)
  if err != nil {
    if os.IsPermission(err) {
      logMessage("[System] ERROR: Insufficient permissions to read /proc/kmsg. Enable CAP_SYSLOG.")
    } else {
      logMessage(fmt.Sprintf("[System] ERROR in IP logger: %v", err))
```

```
    }
    return
  }
  defer f.Close()

  reader := bufio.NewReader(f)
  for {
    line, err := reader.ReadString('\n')
    if err != nil {
      if err == io.EOF {
        time.Sleep(100 * time.Millisecond)
        continue
      }
      logMessage(fmt.Sprintf("[System] ERROR reading kmsg: %v", err))
      time.Sleep(1 * time.Second)
      continue
    }

    if strings.Contains(line, "IP_BLOCK") {
      match := prefixRe.FindStringSubmatch(line)
      if len(match) > 1 {
        prefix := match[1]

        src := "unknown"
        if m := srcRe.FindStringSubmatch(line); len(m) > 1 {
          src = m[1]
        }

        dst := "unknown"
        if m := dstRe.FindStringSubmatch(line); len(m) > 1 {
          dst = m[1]
        }

        proto := "unknown"
        if m := protoRe.FindStringSubmatch(line); len(m) > 1 {
          proto = m[1]
        }

        spt := ""
        if m := sptRe.FindStringSubmatch(line); len(m) > 1 {
          spt = m[1]
        }

        dpt := ""
        if m := dptRe.FindStringSubmatch(line); len(m) > 1 {
          dpt = m[1]
        }

        inIf := ""
        if m := inRe.FindStringSubmatch(line); len(m) > 1 {
          inIf = m[1]
        }

        outIf := ""
        if m := outRe.FindStringSubmatch(line); len(m) > 1 {
          outIf = m[1]
        }

        direction := "inbound"
        if strings.HasSuffix(prefix, "DST") {
          direction = "outbound"
        }

        listType := "MALICIOUS"
        if !strings.Contains(prefix, "MALICIOUS") {
          parts := strings.Split(prefix, "_")
          if len(parts) >= 3 {
            listType = "GEO_" + parts[2]
          }
        }

        portStr := ""
        if dpt != "" {
          portStr = ":" + dpt
        }

        srcPortStr := ""
        if spt != "" {
          srcPortStr = ":" + spt
        }

```

```
        ifStr := ""
        if inIf != "" && outIf != "" {
          ifStr = fmt.Sprintf("IF: %s->%s", inIf, outIf)
        } else if inIf != "" || outIf != "" {
          ifStr = fmt.Sprintf("IF: %s%s", inIf, outIf)
        }

        logMessage(fmt.Sprintf("[IP] Blocked %s to %s%s (Proto: %s) from %s%s (%s) | List: %s", direction, dst,
      portStr, proto, src, srcPortStr, ifStr, listType))
      }
    }
  }
}

func main() {
  logMessage("[System] Nullexit Block Logger starting...")

  go monitorDNS()
  monitorIPs()
}
```

38 scripts/pf.conf

```
# scripts/pf.conf - Packet Filter configuration for nullexit killswitch
# Dynamically parameterized by pfctl macro override

# ext_if must be passed via command line, e.g. -D ext_if=en0
# If not passed, default to en0
ext_if = "en0"
ts_if = "utun+"

# Tables populated dynamically via pfctl
table <vpn_endpoints> persist
table <derp_relays> persist

# Core Rules
set skip on lo0                 # Ensure local loopback is never blocked
scrub out all max-mss 1160      # TCP MSS Clamping: Prevents fragmentation on host TCP flows (like Tailscale
     control plane)
pass quick on $ts_if all        # Allow all Tailscale traffic (essential to prevent SSH lockout)

# Default deny outbound WAN traffic
block out on $ext_if all

# Block inbound Screen Sharing (VNC) and SSH from local Wi-Fi, forcing them to use Tailscale
block in on $ext_if proto tcp from any to any port { 22, 5900 }

# Exceptions for tunnel handshakes and coordination
pass out quick on $ext_if proto udp to <vpn_endpoints> port 2408
pass out quick on $ext_if proto udp to <derp_relays>
pass out quick on $ext_if proto tcp to 192.200.0.0/24 port 443

# Allow local LAN traffic (AirDrop, local Tailscale P2P, etc)
pass out quick on $ext_if proto tcp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
pass out quick on $ext_if proto udp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
pass out quick on $ext_if proto icmp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }

# Allow DHCP outbound (client port 68 to server port 67) so we can obtain/renew IP address
pass out quick on $ext_if proto udp from any port 68 to any port 67
```

39 scripts/reinstall-host-tailscale.sh

```
#!/usr/bin/env bash
#
# scripts/reinstall-host-tailscale.sh
#
# Complete uninstall + reinstall of the host-side Homebrew Tailscale.
#
# Use this when the brew LaunchAgent is wedged, Login Items entries are
# duplicated, "Tailscale is stopped" persists despite
# `brew services restart tailscale`, or you simply want a clean baseline
# before re-engaging the gateway as exit node.
```

```
#
# Idempotent. Safe to re-run. Use --dry to preview every step with zero
# changes. Use --yes to skip the per-step confirmation prompts.
#
# This script ONLY touches the HOST-side Tailscale (the one installed by
# Homebrew `brew install tailscale` and managed by launchd as
# homebrew.mxcl.tailscale). The nullexit gateway container's tailscaled
# (which lives inside the Colima VM at 100.100.21.8 and offers the exit
# node to this host) is unaffected.

set -u

SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"

DRY=0
YES=0

for arg in "$@"; do
  case "$arg" in
    --dry)  DRY=1 ;;
    --yes)  YES=1 ;;
    --help|-h)      # NOTE: heredoc is UNQUOTED so $SCRIPT_DIR expands in the body.  If you
      # add new $-style placeholders here, they'll expand the same way.
      cat <<USAGE
Usage: bash "$SCRIPT_DIR/reinstall-host-tailscale.sh" [--dry] [--yes]

  --dry  Print every step, make zero changes. Use to audit before
         committing.
  --yes  Auto-confirm the prompts at every destructive step. Useful in
         CI / fully-scripted recovery; do not use until you've read
         the script.

After this script completes cleanly, you still must (MANUALLY):

  1. Open System Settings  Privacy & Security  Local Network and
     enable Tailscale. macOS will prompt automatically the first time
     the fresh install of `tailscale` is invoked.

  2. Re-engage the tailnet + exit node:
        sudo tailscale up --ssh=true --accept-dns=false \
          --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
          --exit-node-allow-lan-access=true

  3. Re-run ./toggle.sh from the project root.

  4. Final verification:
        curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
     expected: warp=on, ip=<Cloudflare>, colo=YYZ/YUL
USAGE
      exit 0
      ;;
    *)
      echo "Unknown flag: $arg" >&2
      exit 2
      ;;
  esac
done

# ---------- helpers ----------

confirm() {
  if [ "$YES" = 1 ] || [ "$DRY" = 1 ]; then
    return 0
  fi
  printf "      %s  [y/N] " "$1"
  read -r ans
  case "$ans" in
    y|Y|yes|YES) return 0 ;;
    *)
      echo "     aborted by user"
      exit 1
      ;;
  esac
}

header() {
  printf '\n %s\n' "$1"
}

# with_timeout <seconds> <cmd...>
# Run a command in the background and kill it if it exceeds <seconds>.
```

```
# Returns the command's exit code if it finished in time, 124 (matching
# GNU `timeout` convention) if it timed out. macOS doesn't ship GNU
# coreutils `timeout`, so this is the bash-3.2-friendly replacement.
# stdout + stderr are discarded (we only care about exit code); redirect
# or pipe externally if you need the output.
with_timeout() {
  local timeout_s="${1:-30}"
  shift
  "$@" >/dev/null 2>&1 &
  local cmdpid=$!
  local elapsed=0
  while kill -0 "$cmdpid" 2>/dev/null; do
    if [ "$elapsed" -ge "$timeout_s" ]; then
      kill -9 "$cmdpid" 2>/dev/null
      wait "$cmdpid" 2>/dev/null
      return 124
    fi
    sleep 1
    elapsed=$((elapsed+1))
  done
  wait "$cmdpid" 2>/dev/null
  return $?
}

# ---------- 0. Pre-flight ----------
header "Pre-flight  current host tailscale state"
echo "    brew:                  $(command -v brew >/dev/null 2>&1 && brew --version | head -1 || echo 'not on PATH')"
echo "    tailscaled running:    $(pgrep -lf tailscaled 2>/dev/null | wc -l | tr -d ' ') process(es)"
echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{print $2, $3}' | tr '\n' ' ' | sed 's/ $//')"
echo "    LaunchAgent plists (under ~/Library/LaunchAgents/):"
ls -1 ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/^/      /' || echo "      (none)"

# Detect Tailscale's macOS Network Extension (residual from prior App Store
# Tailscale.app installs). Lives under /Library/Apple/SystemExtensions/ and
# is managed by `systemextensionsctl`  NOT launchd. If loaded, brew
# tailscaled will NOT engage the Network Extension framework slot because
# the SE still has it claimed. Symptom: 'sudo tailscale status' returns
# "Tailscale is stopped" even when the daemon is alive and listening.
se_found=0
if command -v systemextensionsctl >/dev/null 2>&1; then
  # Grab any line matching tailscale and network-extension
  se_line=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale\..*network-extension' | head -n 1)
  if [ -n "$se_line" ]; then
    se_found=1
    # Extract the teamID (10-char uppercase/numeric word)
    se_teamID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^[A-Z0-9]{10}$' | head -n 1)
    # Extract the bundleID (starts with io.tailscale or com.tailscale)
    se_bundleID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '^(io|com)\.tailscale\.' | head -n 1 | sed 's/(.*//')

    # If teamID is not found, fallback to '-'
    if [ -z "$se_teamID" ]; then
      se_teamID="-"
    fi
  fi
fi

if [ "$se_found" = 1 ] && [ -n "$se_bundleID" ]; then
  echo "    Tailscale Network Extension System Extension is loaded (residual from App Store Tailscale.app install)"
  echo "      Team ID: $se_teamID, Bundle ID: $se_bundleID"
  echo "      brew tailscaled cannot engage the NE slot while this SE has it claimed."
  if [ "$DRY" = 1 ]; then
    echo "    [dry] (would prompt) sudo systemextensionsctl uninstall $se_teamID $se_bundleID"
  else
    confirm "Run 'sudo systemextensionsctl uninstall $se_teamID $se_bundleID'? (frees the NE slot for brew tailscale; non-fatal if it errors)"
    sudo systemextensionsctl uninstall "$se_teamID" "$se_bundleID" 2>&1 \
      || echo "    ! uninstall returned non-zero (may require reboot, non-fatal)"
  fi
else
  echo "    no Tailscale System Extension loaded"
fi

# Pattern used by the Step-1 bootout loops, in both dry + real branches.
# Widened from a bare /tailscale/ so the same loops also pick up any
# nullexit-labeled LaunchAgents (e.g., com.nullexit.wake-recovery, the
# caffeinate + post-wake recovery hook). The wake-recovery LaunchAgent
```

```
# is re-installed by ./toggle.sh / setup.sh on re-engage, so wiping it
# here is non-fatal and is part of restoring a clean baseline.
TAILSCALE_LABEL_RE='tailscale'
NULLEXIT_LABEL_RE='nullexit'

# ---------- 1. brew services stop ----------
header "Step 1  Stop the brew-managed service + unload LaunchAgents"
# Detect whether the brew-managed service is registered as $USER (gui/$UID
# domain) or as root (system domain). NOPASSWD sudoers can leave a stale
# root-owned registration after a prior 'sudo brew services ...' round;
# 'brew services stop' without sudo refuses to touch a root-owned plist
# (Error: Service 'tailscale' is started as `root`).
brew_status_line="$(brew services list 2>/dev/null | awk '$1=="tailscale"')"
brew_status_user="$(echo "$brew_status_line" | awk '{print $3}')"
if [ -n "$brew_status_line" ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew services stop tailscale (user=$brew_status_user)"
  else
    if [ "$brew_status_user" = "root" ]; then
      confirm "Run 'sudo brew services stop tailscale' (service is registered as root)?"
      sudo brew services stop tailscale 2>&1 || echo "    (sudo brew services stop returned non-zero  non-fatal
    )"
    else
      echo "     brew services stop tailscale (user=$brew_status_user, no sudo needed)"
      brew services stop tailscale 2>&1 || echo "    (brew reported it wasn't actually running)"
    fi
  fi
  echo "     unload any tailscale-related LaunchAgents (prevents launchd from respawning after kill)"
  if [ "$DRY" = 1 ]; then
    echo "      [dry] for each 'tailscale|nullexit'-matched label in gui/$UID domain:"
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "            [dry] launchctl bootout gui/$UID/$label"
    done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 &&
      ($3 ~ ts || $3 ~ nx) {print $3}')
    echo "      [dry] for each 'tailscale|nullexit'-matched label in system domain (root-owned brew plist):"
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "            [dry] (would prompt) sudo launchctl bootout system/$label"
    done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
     NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
    if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
      echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macos"
    fi
    if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
      echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macsys"
    fi
  else
    # Bootout every gui-domain LaunchAgent whose label matches 'tailscale'
    # OR 'nullexit'. The nullexit arm catches com.nullexit.wake-recovery,
    # which holds caffeinate/watcher.sh alive across a daemon reset  not
    # desired while the script is establishing a fresh baseline. Awake-
    # recovery is re-installed by ./toggle.sh / setup.sh on re-engage, so
    # wiping here is non-fatal. The pattern stays narrower than a bare
    # /tail/ or /null/ so unrelated Apple services are not unloaded.
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      echo "       launchctl bootout gui/$UID/$label"
      launchctl bootout "gui/$UID/$label" 2>/dev/null || true
    done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 &&
      ($3 ~ ts || $3 ~ nx) {print $3}')
    # Bootout every system-domain LaunchDaemon/Agent whose label matches
    # 'tailscale' OR 'nullexit'. This catches the root-owned brew plist left
    # behind by prior 'sudo brew services ...' invocations  that was the
    # case that left tailscaled PID 47447  57734 respawning between sudo
    # pkill runs. Defensive nullexit catch in case the user previously
    # `sudo ln -s`'d the wake-recovery plist into /Library/LaunchDaemons.
    while IFS= read -r label; do
      [ -z "$label" ] && continue
      confirm "Run 'sudo launchctl bootout system/$label'?"
      sudo launchctl bootout "system/$label" 2>/dev/null || true
    done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
     NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
    # Legacy Tailscale.app system-domain entries. Each one is checked and
    # confirmed separately so we don't sudo-promp or sudo-bootout for nothing.
    if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
      confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macos'?"
      sudo launchctl bootout system/com.tailscale.ipn.macos 2>/dev/null || true
    fi
    if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
      confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macsys'?"
```

```
      sudo launchctl bootout system/com.tailscale.ipn.macsys 2>/dev/null || true
    fi
  fi
  echo "     waiting up to 10s for tailscaled to exit"
  for i in 1 2 3 4 5 6 7 8 9 10; do
    if [ "$DRY" = 1 ]; then
      echo "    [dry] sleep 1; loop check would terminate here"
      break
    fi
    sleep 1
    if ! pgrep -lf tailscaled >/dev/null 2>&1; then
      echo "     tailscaled exited in ${i}s"
      break
    fi
    [ "$i" = 10 ] && echo "     tailscaled still alive after 10s  Step 2 will deal with it"
  done
else
  echo "     brew reports no tailscale service loaded"
fi

# ---------- 2. Kill any orphan tailscaled ----------
header "Step 2  Kill any orphan tailscaled process(es) (may escalate to sudo)"
if ! pgrep -lf tailscaled >/dev/null 2>&1; then
  echo "     no orphans"
else
  echo "    Stray process(es):"
  pgrep -lf tailscaled | sed 's/^/      /'
  echo "     no-sudo: pkill tailscaled"
  if [ "$DRY" = 1 ]; then
    echo "    [dry] pkill tailscaled  on failure: sudo pkill -9 tailscaled"
  else
    if pkill tailscaled 2>/dev/null; then
      echo "     no-sudo pkill succeeded"
    else
      echo "    ! no-sudo pkill insufficient; cannot escalate yet"
    fi
    sleep 1
  fi

  # Recheck; if still alive, offer sudo escalation
  if { [ "$DRY" = 0 ] && pgrep -lf tailscaled >/dev/null 2>&1; } || [ "$DRY" = 1 ]; then
    if [ "$DRY" = 1 ]; then
      echo "    [dry] sudo pkill -9 tailscaled  (escalation would fire here)"
    else
      confirm "Run 'sudo pkill -9 tailscaled' to escalate?"
      sudo pkill -9 tailscaled 2>&1
      sleep 1
    fi
  fi

  if [ "$DRY" = 0 ]; then
    if pgrep -lf tailscaled >/dev/null 2>&1; then
      echo "     Failed to kill stragglers  bailing out"
      pgrep -lf tailscaled | sed 's/^/      /'
      exit 3
    fi
    echo "     all tailscaled processes gone"
  fi
fi

# ---------- 3. brew uninstall ----------
header "Step 3  brew uninstall tailscale"
if ! brew list tailscale >/dev/null 2>&1; then
  echo "     not currently installed via brew  skip"
else
  # Detect whether the keg directory is owner=$USER or owner=root. If the
  # keg is root-owned (typically a residual of some prior 'sudo brew ...'
  # cycle that did perms fixups), 'brew uninstall' without sudo refuses
  # to remove the files and prints: 'Error: Could not remove tailscale
  # keg! Do so manually: sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>'.
  # Detect that and escalate to a sudoed uninstall, with a manual rm -rf
  # fallback if `sudo brew uninstall` itself errors out (e.g., partial
  # state where the keg dir is gone but brew still thinks it's installed).
  # Detect keg state. Treat `(dir absent)` AND `(dir present + empty)` as the
  # same `<unknown>` bucket so both escalate to sudo. The empty-parent case
  # is what you get after a manual `sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>`
  # from inside: brew's parent dir is left present-but-empty and brew still
  # thinks the package is installed.
  keg_dir="/opt/homebrew/Cellar/tailscale"
  if [ -d "$keg_dir" ] && [ -n "$(ls -A "$keg_dir" 2>/dev/null)" ]; then
    keg_owner="$(stat -f '%Su' "$keg_dir" 2>/dev/null || echo '<unknown>')"
```

```
  else
    keg_owner="<unknown>"
  fi
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew uninstall tailscale (keg_owner=$keg_owner)"
  else
    # 'root' AND '<unknown>' both require sudo escalation. The '<unknown>'
    # case is the partial-state path: keg dir was already removed (manually
    # or by a prior failed uninstall) but brew's internal state still says
    # installed. Without this branch the script would fall through to a
    # non-sudo brew uninstall that fails again with the same 'Could not
    # remove tailscale keg!' error.
    if [ "$keg_owner" = "root" ] || [ "$keg_owner" = "<unknown>" ]; then
      confirm "Run 'sudo brew uninstall tailscale' (keg=$keg_owner)?"
      sudo brew uninstall tailscale 2>&1 || {
        echo "    ! sudo brew uninstall failed; offering manual rm -rf fallback"
        # Guard the glob so an empty/absent dir doesn't literal-expand the
        # '*' and confuse the user with a sudo error about a non-existent path.
        if [ -d /opt/homebrew/Cellar/tailscale ] && [ -n "$(ls -A /opt/homebrew/Cellar/tailscale 2>/dev/null)"
    ]; then
          confirm "Run 'sudo rm -rf /opt/homebrew/Cellar/tailscale/*'?"
          sudo rm -rf /opt/homebrew/Cellar/tailscale/* 2>&1
        else
          echo "     keg dir already empty or absent  nothing more to rm"
        fi
      }
    else
      confirm "Run 'brew uninstall tailscale' (removes LaunchAgent plist + linked symlinks)?"
      brew uninstall tailscale 2>&1
    fi
  fi
fi

# ---------- 4. Wipe state directories ----------
header "Step 4  Wipe stale state"
# User-owned
declare -a USER_PATHS
USER_PATHS=(
  "$HOME/.local/share/tailscale"
  "$HOME/.local/run/tailscaled.socket"
  "$HOME/.local/run/tailscaled.state"
  "$HOME/.local/state/tailscale"
  "$HOME/Library/Application Support/Tailscale"
  "$HOME/Library/Caches/com.tailscale.ipn.macos"
  "$HOME/Library/Caches/Tailscale"
  "$HOME/Library/Logs/Tailscale"
)
echo "    User-owned paths (no sudo):"
for p in "${USER_PATHS[@]}"; do
  if [ -e "$p" ]; then
    if [ "$DRY" = 1 ]; then
      echo "      [dry] rm -rf '$p'"
    else
      rm -rf "$p"
      echo "       rm -rf '$p'"
    fi
  fi
done

# System
declare -a SUDO_PATHS
SUDO_PATHS=(
  "/var/lib/tailscale"
  "/var/cache/tailscale"
)
echo "    System paths under /var (will prompt for sudo if any are present):"
eligible_sudo=0
for p in "${SUDO_PATHS[@]}"; do
  if [ -e "$p" ]; then
    eligible_sudo=1
    if [ "$DRY" = 1 ]; then
      echo "      [dry] sudo rm -rf '$p'"
    else
      if confirm "Run 'sudo rm -rf $p'?"; then
        sudo rm -rf "$p" && echo "       rm -rf '$p'"
      fi
    fi
  fi
done
[ "$eligible_sudo" = 0 ] && echo "      (none of /var/lib/tailscale, /var/cache/tailscale are present)"

```

```
# ---------- 5. Wipe stale LaunchAgent plists ----------
header "Step 5  Wipe stale LaunchAgent plist files"
declare -a PLIST_FILES
PLIST_FILES=(
  "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
  "$HOME/Library/LaunchAgents/com.tailscale.ipn.macos.plist"
  "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist.bak"
  "$HOME/Library/LaunchAgents/com.nullexit.wake-recovery.plist"
)
for f in "${PLIST_FILES[@]}"; do
  if [ -e "$f" ]; then
    if [ "$DRY" = 1 ]; then
      echo "    [dry] rm -f '$f'"
    else
      rm -f "$f"
      echo "     rm -f '$f'"
    fi
  fi
done
echo "    Remaining tailscale|nullexit plists (should be empty):"
ls -1 ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/^/      /' || echo "      (none)"

# ---------- 5.5. Drain Background-Items (Login Items) ----------
header "Step 5.5  Drain Background-Items (Login Items)"
# macOS 13+ keeps a parallel "Background Items" database alongside
# classic LaunchAgents. Even when we wipe ~/Library/LaunchAgents/homebrew.mxcl.tailscale.plist
# in Step 5, the .btm file at
# ~/Library/Application Support/com.apple.backgroundtaskmanagementagent.backgrounditems.btm
# can still hold a Background-Items entry pointing at the (now-dead) plist
# path. Next login: macOS sees the entry, tries to bootstrap, finds the
# plist missing, and either silently regenerates it via brew's post-install
# hook (causing tailscaled respawn at next login) or refuses outright.
# `sfltool reset-login-items` is Apple's canonical API for wiping the
# entire .btm file. It DOES wipe all Login Items  the user re-adds what
# they want via System Settings  General  Login Items. Acceptable here
# because the goal of this script is "fresh baseline for Tailscale".
BTM="$HOME/Library/Application Support/com.apple.backgroundtaskmanagementagent/backgrounditems.btm"
if [ -f "$BTM" ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] found $BTM; would prompt sfltool reset-login-items"
  else
    confirm "Run 'sfltool reset-login-items' (wipes ALL Login Items; rebuild via System Settings  General
     Login Items)?"
    if sfltool reset-login-items 2>&1; then
      echo "     Login Items drained"
    else
      echo "    ! sfltool reset-login-items failed (non-fatal; do manually: System Settings  General  Login
     Items)"
    fi
  fi
else
  echo "     no backgrounditems.btm present"
fi

# ---------- 6. brew install ----------
header "Step 6  brew install tailscale (fresh)"
if brew list tailscale >/dev/null 2>&1; then
  echo "     already installed  skip"
else
  confirm "Run 'brew install tailscale'?"
  if [ "$DRY" = 1 ]; then
    echo "    [dry] brew install tailscale"
  else
    brew install tailscale 2>&1
  fi
fi

# ---------- 6.5. Strip quarantine xattr from keg ----------
header "Step 6.5  Strip quarantine xattr from tailscale install"
# brew's freshly poured bottles don't carry com.apple.quarantine (the
# bottles are content-addressed and Homebrew-verified), but the upstream
# installer (Tailscale.app from App Store) and any manual download do.
# Strip the xattr from the entire keg as a defensive safety net so the
# freshly-installed tailscaled binary and any helper binaries are not
# Gatekeeper-blocked on first launch. Idempotent  no-op if no quarantine
# xattr is present.
quarantine_count="$(xattr -lr /opt/homebrew/Cellar/tailscale 2>/dev/null | grep -c '^[^ ]* com.apple.quarantine
    ' || echo 0)"
if [ "${quarantine_count:-0}" -gt 0 ]; then
  if [ "$DRY" = 1 ]; then
    echo "    [dry] found $quarantine_count file(s) with com.apple.quarantine; would prompt xattr -dr"
```

```
  else
    confirm "Run 'xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale'? (strips quarantine xattr from $quarantine_count file(s); non-fatal if it errors)"
    if xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale 2>&1; then
      echo "     quarantine xattr stripped"
    else
      echo "    ! xattr -dr failed (non-fatal; open /opt/homebrew/Cellar/tailscale in Finder, right-click tailscaled  Open Anyway)"
    fi
  fi
else
  echo "     no com.apple.quarantine xattr on tailscale install"
fi

# ---------- 7. brew services start (NO sudo!) ----------
header "Step 7  Start the service as $USER (NO sudo, with multi-tier self-healing)"
if ! command -v tailscale >/dev/null 2>&1; then
  echo "     tailscale not on PATH after install  brew install errored; aborting"
  exit 4
fi

# Aliveness fan-out: ANY one of these is taken as proof of life. None of
# them alone is reliable  pgrep misses tailscaled if it does setproctitle
# or fork-detach; launchctl's reported PID can be a shim that exited
# within 1s; tailscale-cli hangs while macOS Local-Network permission is
# unresolved. Combined, they cover each other's blind spots.

tailscaled_api_up() {
  # Strategy A: tailscale CLI can talk to the local API. Canonical proof.
  # Bound at 8s  tailscale status itself returns quickly on success OR
  # failure, but if the post-install state is wedged it can hang.
  with_timeout 8 /opt/homebrew/bin/tailscale status >/dev/null 2>&1
}

tailscaled_proc_up() {
  # Strategy B: any process has 'tailscaled' substring in argv.
  pgrep -lf tailscaled >/dev/null 2>&1
}

tailscaled_lc_up() {
  # Strategy C: launchctl-managed service has a non-dash, non-zero PID.
  local lc_pid
  lc_pid="$(launchctl list 2>/dev/null | awk '$3 == "homebrew.mxcl.tailscale" {print $1; exit}')"
  if [ -n "$lc_pid" ] && [ "$lc_pid" != "-" ] && [ "$lc_pid" -gt 0 ] 2>/dev/null; then
    return 0
  fi
  return 1
}

tailscaled_is_alive() {
  tailscaled_api_up  && return 0
  tailscaled_proc_up && return 0
  tailscaled_lc_up   && return 0
  return 1
}

confirm "Run 'brew services start tailscale'?"
if [ "$DRY" = 1 ]; then
  echo "    [dry] brew services start tailscale"
  echo "    [dry] (then poll up to 30s for tailscaled_is_alive)"
  echo "    [dry] (recovery tier 1: launchctl kickstart)"
  echo "    [dry] (recovery tier 2: launchctl bootout + bootstrap of explicit plist)"
  echo "    [dry] (recovery tier 3: brew services restart (stop+start cycle))"
  echo "    [dry] (recovery tier 4: direct /opt/homebrew/bin/tailscaled spawn (bypasses launchd))"
else
  brew services start tailscale 2>&1
  # Primary poll  tailscaled_is_alive covers all 3 detection strategies.
  echo "     polling for tailscaled to come up (up to ~30s)"
  tailscaled_up=0
  for i in 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15; do
    if tailscaled_is_alive; then
      tailscaled_up=1
      echo "       tailscaled alive after ${i}*2s"
      break
    fi
    sleep 2
  done

  # Recovery tier 1: kickstart the launchd-managed service. This is the
  # cheapest and most-aligned-with-launchd fix when brew has loaded the
  # LaunchAgent but tailscaled didn't fork. Kickstart tells launchd
```

```
  # "re-spawn if not running" without forcing a teardown.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 1: launchctl kickstart"
    if launchctl kickstart -kp "gui/$UID/homebrew.mxcl.tailscale" 2>&1; then
      for i in 1 2 3 4 5 6 7 8 9 10; do
        if tailscaled_is_alive; then
          tailscaled_up=1
          echo "       tailscaled alive after kickstart (${i}*2s)"
          break
        fi
        sleep 2
      done
    else
      echo "      ! kickstart returned non-zero (LaunchAgent may not be loaded)"
    fi
  fi

  # Recovery tier 2: explicit plist bootout + bootstrap. Forces a fresh
  # launchd registration cycle for the brew-managed service, which fixes
  # the wider class of "brew says started but daemon never spawned"
  # wedge states that kickstart alone can't dig out of.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 2: launchctl bootout + bootstrap of explicit plist"
    launchctl bootout "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null || true
    sleep 1
    # Brew installs the canonical plist template at $HOMEBREW_PREFIX/Library/.
    # On Apple Silicon that's /opt/homebrew/Library/LaunchAgents/. On Intel
    # it's /usr/local/Library/LaunchAgents/. Detect which and pick the
    # right one so this script works across both architectures.
    if [ -f /opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
      plist_path="/opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
    elif [ -f /usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
      plist_path="/usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
    else
      plist_path=""
    fi
    if [ -n "$plist_path" ]; then
      if launchctl bootstrap "gui/$UID" "$plist_path" 2>&1; then
        for i in 1 2 3 4 5 6 7 8 9 10; do
          if tailscaled_is_alive; then
            tailscaled_up=1
            echo "       tailscaled alive after bootstrap (${i}*2s)"
            break
          fi
          sleep 2
        done
      else
        echo "      ! explicit bootstrap returned non-zero"
      fi
    else
      echo "      ! could not locate brew plist template; skipping bootstrap tier"
    fi
  fi

  # Recovery tier 3: brew services restart (forces a stop + start cycle,
  # which unsticks LaunchAgents with stale 'started' state from prior
  # half-cycles). This is the canonical brew-doc fix for 'services don't
  # actually start even though brew says started'.
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 3: brew services restart (forces stop+start cycle)"
    if brew services restart tailscale 2>&1; then
      for i in 1 2 3 4 5 6 7 8 9 10; do
        if tailscaled_is_alive; then
          tailscaled_up=1
          echo "       tailscaled alive after brew restart (${i}*2s)"
          break
        fi
        sleep 2
      done
    else
      echo "      ! brew services restart returned non-zero"
    fi
  fi

  # Recovery tier 4: last-resort direct spawn of the daemon binary,
  # bypassing launchd entirely. If tailscaled crashes here too, we'll
  # capture its actual stderr in /tmp for the user to read (NOPASSWD
  # sudoers often have local-network permission denied  the captured
  # stderr makes it obvious what to fix in System Settings).
  if [ "$tailscaled_up" = 0 ]; then
    echo "    ! Recovery tier 4: direct spawn (bypasses launchd; tail stderr to /tmp)"
```

```
    /opt/homebrew/bin/tailscaled >/tmp/nullexit-tailscaled-spawn.log 2>&1 &
    spawned_pid=$!
    disown "$spawned_pid" 2>/dev/null || true
    sleep 3
    if tailscaled_is_alive; then
      tailscaled_up=1
      echo "       tailscaled alive via direct spawn (pid=$spawned_pid)"
      echo "      (NOTE: not under launchd supervision  a reboot, re-run of"
      echo "       this script, or ./toggle.sh is required to recreate."
      echo "       Last 5 lines of stderr captured to /tmp/nullexit-tailscaled-spawn.log:)"
      tail -n 5 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/        /' || true
    fi
  fi

  if [ "$tailscaled_up" = 0 ]; then
    echo "     HARD FAIL: tailscaled refuses to stay alive across every recovery tier"
    echo "      launchctl-managed state (if available):"
    launchctl print "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null \
      | grep -E 'state =|last exit code =|runs =|path =' \
      | sed 's/^/        /' \
      || echo "        (no launchctl state available for homebrew.mxcl.tailscale)"
    echo "      Last 10 lines of /tmp/nullexit-tailscaled-spawn.log (if any):"
    tail -n 10 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/^/        /' || echo "        (no log)"
    echo "      Most likely remaining cause: macOS Local Network permission denied."
    echo "      System Settings  Privacy & Security  Local Network  toggle Tailscale ON."
    exit 7
  fi
fi
sleep 3

# ---------- 8. Verify ----------
header "Step 8  Verify fresh install"
echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{print $2, $3}' | tr '\n' ' ' | sed 's/ $//')"
echo "    tailscaled process:     $(pgrep -lf tailscaled 2>/dev/null | head -1 || echo 'NONE (check brew logs)')"
ts_ver=""
if [ -x /opt/homebrew/bin/tailscale ]; then
  ts_ver="$(/opt/homebrew/bin/tailscale version 2>/dev/null | head -1 | tr -d '\n')"
elif command -v tailscale >/dev/null 2>&1; then
  ts_ver="$(tailscale version 2>/dev/null | head -1 | tr -d '\n')"
fi
echo "    tailscale version:      ${ts_ver:-binary missing}"
ts_status=""
if [ -x /opt/homebrew/bin/tailscale ]; then
  ts_status="$(/opt/homebrew/bin/tailscale status 2>/dev/null | head -3)"
fi
if [ -n "$ts_status" ]; then
  echo "    tailscale status:"
  printf '%s\n' "$ts_status" | sed 's/^/      /'
else
  echo "    tailscale status:       (binary missing or errored)"
fi

if { [ "$DRY" = 0 ] && /opt/homebrew/bin/tailscale status 2>/dev/null | grep -q '^Tailscale is stopped'; }; then
  echo
  echo "     NOTICE: 'tailscale status' still reports stopped after a fresh install."
  echo "       Two equally likely causes:"
  echo "         (a) macOS hasn't yet shown the first-run 'Local Network' prompt for"
  echo "             Tailscale  open System Settings  Privacy & Security  Local"
  echo "             Network. Toggle Tailscale ON."
  echo "         (b) macOS denied Local Network permission at some prior point."
  echo "       Either way: System Settings  Privacy & Security  Local Network."
  echo "       If the entry isn't present, invoking any tailscale command will"
  echo "       re-trigger the prompt automatically."
fi

# ---------- 9. Done ----------
header "Done"
if [ "$DRY" = 1 ]; then
  echo "    (dry-mode finished; no follow-up commands were issued.)"
  exit 0
fi

cat <<'DONE'

   Manual follow-ups (all of these; the script stops short of running sudo
    tailscale up so you can verify the brew install is actually live first):

    1.  System Settings  Privacy & Security  Local Network  Tailscale ON
```

```
       macOS will normally pop the prompt the first time you run `tailscale`.
       If it didn't, trigger it via the menu bar  check.

    2.  Re-engage the tailnet + exit node:
           sudo tailscale up --ssh=true --accept-dns=false \
             --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
             --exit-node-allow-lan-access=true

    3.  Re-run: ./toggle.sh
       (re-installs the com.nullexit.wake-recovery LaunchAgent  wiped
       in Step 5 to clear the slate  and re-engages the gateway)

    4.  Final verification:
           curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
         expect: warp=on, ip=<Cloudflare IP>, colo=YYZ/YUL

    Container-side (gateway / 100.100.21.8) Tailscale is unaffected by this
    script. Other tailnet devices (S24, iPhone, Windows) are unaffected.

DONE
```

40 scripts/routing-fix.sh

```
#!/bin/sh
# routing-fix: Maintains routing for SOCKS5 proxy traffic through WARP (tun0)
#              and FORWARD rules for Tailscale exit-node return traffic.
#
# Runs inside the warp container's network namespace. This script ensures:
# 1. Table 200 has default via tun0 (for SOCKS5 proxy traffic from 172.18.0.2)
#    with the WARP endpoint exception via eth0 (prevents tunnel loop)
# 2. The CGNAT ip rule (100.64.0.0/10  table 52  tailscale routes) is maintained for Tailscale
# 3. FORWARD RELATED,ESTABLISHED rule allows return traffic for exit-node
#    forwarded connections (S24  tailscale0  eth0  host  internet)
#
# CRITICAL: Does NOT touch the main table's default route. Gluetun handles that
# internally via its own routing rules and WireGuard fwmark (0xca6c  table 51820).
# Overriding the main table default would create a loop: encrypted WireGuard
# packets exiting tun0 would match default dev tun0 and re-enter the tunnel.

set -e

trap 'echo "routing-fix: Caught signal, exiting..."; exit 0' TERM INT QUIT

sleep 15 &
wait $! || true

echo "routing-fix: Setting up routing for SOCKS5 proxy through WARP (tun0)..."

# Dynamically detect Docker subnet from eth0
DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')
DOCKER_GW=$(ip route show default | head -1 | awk '{print $3}')
WARP_ENDPOINT="${WARP_ENDPOINT}"
TAILSCALE_ALLOW_LAN_P2P="${TAILSCALE_ALLOW_LAN_P2P:-false}"
IP_BLOCKLIST_FILE="/userfilters/ip_blocklist.ipset"
LAST_IP_MTIME=""
echo "routing-fix: Docker network: ${DOCKER_NET}, gateway: ${DOCKER_GW}"

# Table 200: Used by ip rule 100 (from 172.18.0.2) for SOCKS5 proxy traffic
# WARP endpoint goes via eth0 to avoid tunnel loop, everything else via tun0
ip route add "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || \
  ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
ip route replace default dev tun0 table 200 2>/dev/null || true

# Main table: Just ensure the Docker subnet route exists via eth0
# (gluetun owns the default route here  do NOT touch it)
ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true

# FORWARD RELATED,ESTABLISHED: Allow return traffic for exit-node connections.
# The container has BOTH iptables (nftables backend) and iptables-legacy
# loaded. Gluetun's post-rules.txt uses the nftables backend, while Tailscale
# uses the legacy backend. Packets go through BOTH stacks, so we add the rule
# to both to ensure it survives regardless of which backend has policy DROP.
#
# Packet flow: S24 outgoing (tailscale0->eth0) ACCEPTed by first rule.
# Return traffic (eth0->eth0 via table 52 to host) needs RELATED,ESTABLISHED
# to match the conntrack entry created by the outgoing flow.
```

```
add_fwd_related_established() {
  # nftables backend (used by gluetun's post-rules.txt)
  if ! iptables -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
    iptables -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
  fi
  # legacy backend (used by Tailscale's ts-forward)
  if command -v iptables-legacy >/dev/null 2>&1; then
    if ! iptables-legacy -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
      iptables-legacy -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
    fi
  fi
}

add_fwd_related_established

# Policy Routing for Tailscale P2P:
# Tailscale sends its encrypted mesh UDP packets from port 41642.
# If these go out tun0 (WARP), Cloudflare's strict NAT drops the returning
# hole-punch packets, causing Tailscale to fail P2P and fall back to DERP relays (high latency).
# We mark packets originating from port 41642 and force them out the main table (eth0).
#
# LAN P2P is controlled by two sources (in priority order):
#   1. /app/.lan_p2p_detected  written by watcher.sh on macOS on every network change
#      using system_profiler SPAirPortDataType (802.1x detection) + AP isolation probe. Overrides .env.
#   2. TAILSCALE_ALLOW_LAN_P2P env var (from .env)  static fallback.
# When disabled, we explicitly DROP local RFC1918-destined Tailscale UDP so the
# phone cleanly falls back to DERP rather than getting poisoned by macOS gvproxy SNAT.
add_tailscale_p2p_bypass() {
  # Dynamic override from watcher.sh (network auto-detection) takes priority over .env
  local allow_p2p="${TAILSCALE_ALLOW_LAN_P2P:-false}"
  if [ -f "/app/.lan_p2p_detected" ]; then
    allow_p2p=$(cat "/app/.lan_p2p_detected" 2>/dev/null | tr -d '[:space:]')
  fi

  # Always allow P2P directly to the Docker host gateway (the Mac running this container).
  # The gateway container is always co-located with the host Mac on the same machine,
  # so direct Tailscale P2P via the Docker bridge is always safe and avoids DERP overhead.
  # This RETURN rule must be inserted BEFORE the general DROP rules so it takes priority.
  local default_gw
  default_gw=$(ip route show default | awk '{print $3}')
  if [ -n "$default_gw" ]; then
    if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null; then
      iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null || true
    fi
  fi

  # Also explicitly whitelist all physical IPs of the host Mac (e.g. 172.17.52.223).
  # The Tailscale control plane registers the Mac using its physical IP, so the
  # container will attempt P2P handshakes using that physical IP instead of the
  # Docker bridge IP. If this physical IP falls within 172.16.0.0/12, it would be dropped.
  if [ -f "/app/.host_ips" ]; then
    while IFS= read -r host_ip; do
      [ -z "$host_ip" ] && continue
      if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null; then
        iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null || true
      fi
    done < "/app/.host_ips"
  fi

  if [ "${allow_p2p}" != "true" ]; then
    # Drop Tailscale UDP packets destined for local subnets to prevent macOS gvproxy SNAT
    # endpoint poisoning. When disabled (campus/enterprise WPA2-Enterprise with AP isolation),
    # dropping these forces the S24 to use DERP relays reliably.
    for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
      if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; then
        iptables -t mangle -A OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null || true
      fi
    done
  else
    # Remove DROP rules if they exist (switching from restricted to allowed)
    for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
      while iptables -t mangle -D OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; do :; done
    done
  fi

  # Bypass STUN and P2P packets to public IPs (DERP relays and remote peers) out eth0
  if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null; then
    iptables -t mangle -A OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null || true
  fi
  if ! ip rule show | grep -q 'fwmark 0x8888'; then
    ip rule add fwmark 0x8888 lookup 254 pref 98 2>/dev/null || true
```

```
  fi
}

add_tailscale_p2p_bypass

add_onion_transparent_proxy() {
  for ipt in iptables iptables-legacy; do
    command -v "$ipt" >/dev/null 2>&1 || continue
    if ! $ipt -t nat -C PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null; then
      $ipt -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null || true
    fi
    if ! $ipt -C FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null; then
      $ipt -I FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null || true
    fi
  done
}

add_tailscale_p2p_bypass
add_onion_transparent_proxy

create_log_drop_chain() {
  local ipt="$1"
  local chain="$2"
  local prefix="$3"

  if ! $ipt -N "$chain" 2>/dev/null; then
    $ipt -F "$chain" 2>/dev/null || true
  fi
  $ipt -A "$chain" -j LOG --log-prefix "$prefix"
  $ipt -A "$chain" -j DROP 2>/dev/null || true
}

add_country_block() {
  # To add a country to the blocklist, simply define the 'BLOCKED_COUNTRIES' variable in .env with 2-letter ISO country codes.
  # You can find the country code by looking at the filename on http://www.ipdeny.com/ipblocks/data/countries/
  # For example, to block China (cn) and Russia (ru), set: BLOCKED_COUNTRIES="cn ru" in your .env file
  for zone in $BLOCKED_COUNTRIES; do
    set_name="block_${zone}"
    zone_upper=$(echo "$zone" | tr 'a-z' 'A-Z')
    chain_name="LOG_DROP_${zone_upper}"
    log_prefix="IP_BLOCK_${zone_upper}: "

    # Create the ipset if it doesn't exist
    if ! ipset list "$set_name" >/dev/null 2>&1; then
      ipset create "$set_name" hash:net 2>/dev/null || true
      echo "routing-fix: Downloading ${zone_upper} IPs for blocklist..."
      curl -sS "http://www.ipdeny.com/ipblocks/data/countries/${zone}.zone" | grep -v '^#' | while read -r subnet; do
        [ -n "$subnet" ] && ipset add "$set_name" "$subnet" 2>/dev/null || true
      done
    fi

    # Apply LOG_DROP rule to both backends
    for ipt in iptables iptables-legacy; do
      command -v "$ipt" >/dev/null 2>&1 || continue

      create_log_drop_chain "$ipt" "$chain_name" "$log_prefix"

      # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
      while $ipt -D FORWARD -m set --match-set "$set_name" dst -j DROP 2>/dev/null; do :; done

      if ! $ipt -C FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null; then
        $ipt -I FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null || true
      fi
    done
  done
}

add_ip_blocklist() {
  # Only reload when the compiled file has actually changed (mtime check).
  # This prevents a pointless ipset restore on every 30-second loop tick.
  if [ ! -f "$IP_BLOCKLIST_FILE" ]; then
    return 0
  fi

  CURRENT_MTIME=$(stat -c %Y "$IP_BLOCKLIST_FILE" 2>/dev/null || echo "0")
  if [ "$CURRENT_MTIME" != "$LAST_IP_MTIME" ]; then
    echo "routing-fix: IP blocklist changed, reloading..."
    # ipset restore handles the atomic swap internally:
    # create_new  populate  swap with live  destroy_new
```

```
    if ipset restore < "$IP_BLOCKLIST_FILE" 2>/dev/null; then
      LAST_IP_MTIME="$CURRENT_MTIME"
      echo "routing-fix: IP blocklist loaded ($(ipset list block_malicious | grep -c '^[0-9]') entries)."
    else
      echo "routing-fix: Warning: ipset restore failed. Retrying next cycle."
      return 1
    fi
  fi

  # Apply FORWARD rules in BOTH iptables backends.
  for ipt in iptables iptables-legacy; do
    command -v "$ipt" >/dev/null 2>&1 || continue

    create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_DST "IP_BLOCK_MALICIOUS_DST: "
    create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_SRC "IP_BLOCK_MALICIOUS_SRC: "

    # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
    while $ipt -D FORWARD -m set --match-set block_malicious dst -j DROP 2>/dev/null; do :; done
    while $ipt -D FORWARD -m set --match-set block_malicious src -j DROP 2>/dev/null; do :; done

    if ! $ipt -C FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null; then
      $ipt -I FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null || true
    fi

    if ! $ipt -C FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null; then
      $ipt -I FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null || true
    fi
  done
}

# Start background block logger (DNS + IP)
if ! pgrep -f "nullexit-logger" >/dev/null; then
  echo "routing-fix: Starting background block logger..."
  nullexit-logger &
fi

add_country_block
add_ip_blocklist

echo 'routing-fix: Routes applied.'

UNHEALTHY_COUNT=0

# Re-assert loop (every 30 seconds)
while true; do
  # Table 200 routes
  ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true

  # Tunnel health check
  if ! curl -m 3 -sS --interface tun0 https://cloudflare.com/cdn-cgi/trace >/dev/null 2>&1; then
    UNHEALTHY_COUNT=$((UNHEALTHY_COUNT + 1))
  else
    UNHEALTHY_COUNT=0
    if [ -f "/app/TUNNEL_FAILED_CLOSED.marker" ]; then
      rm -f "/app/TUNNEL_FAILED_CLOSED.marker"
    fi
  fi

  if [ "$UNHEALTHY_COUNT" -ge 3 ]; then
    echo "routing-fix: Tunnel health check failed $UNHEALTHY_COUNT times. Failing closed."
    ip route del default dev tun0 table 200 2>/dev/null || true
    ip route replace blackhole default table 200 2>/dev/null || true
    touch "/app/TUNNEL_FAILED_CLOSED.marker"
  else
    ip route replace default dev tun0 table 200 2>/dev/null || true
  fi

  # Main table: keep Docker subnet route
  ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true

  # CGNAT ip rule for Tailscale
  if ! ip rule show | grep -q '100.64.0.0/10'; then
    ip rule add to 100.64.0.0/10 lookup 52 pref 99 2>/dev/null || true
    echo 'routing-fix: CGNAT rule missing, re-injected'
  fi

  # FORWARD RELATED,ESTABLISHED: Allow return traffic in both iptables
  # backends. Gluetun may reset the nftables ruleset on VPN reconnect,
  # and Tailscale may reset the legacy ruleset on auth refresh.
  add_fwd_related_established

```

```
  # Ensure P2P packets bypass WARP to maintain low latency
  add_tailscale_p2p_bypass

  # Ensure transparent proxying for .onion is active
  add_onion_transparent_proxy

  # Enforce country blocklist
  add_country_block

  # Enforce IP blocklist (reloads automatically when file changes)
  add_ip_blocklist

  sleep 30 &
  wait $! || true
done
```

41 scripts/rule-compiler/go.mod

```
module nullexit/rule-compiler

go 1.22
```

42 scripts/rule-compiler/main.go

```
package main

import (
  "bufio"
  "bytes"
  "crypto/md5"
  "encoding/base64"
  "encoding/hex"
  "fmt"
  "io"
  "net/http"
  "net/netip"
  "os"
  "os/exec"
  "path/filepath"
  "regexp"
  "sort"
  "strings"
  "sync"
  "time"
)

const (
  BlackListPath = "black_list.txt"
  WhiteListPath = "white_list.txt"
  OutputPath    = "adguard/work/userfilters/compiled_rules.txt"
  IpOutputPath  = "adguard/work/userfilters/ip_blocklist.ipset"
  IpCacheDir    = "adguard/work/userfilters/cache/ip"
  CacheDir      = "adguard/work/userfilters/cache"
)

var CoreLists = []string{
  "https://raw.githubusercontent.com/anudeepND/blacklist/master/adservers.txt",
  "https://raw.githubusercontent.com/anudeepND/blacklist/master/facebook.txt",
  "https://raw.githubusercontent.com/crazy-max/WindowsSpyBlocker/master/data/hosts/spy.txt",
  "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.samsung.txt",
  "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.apple.txt",
  "https://abp.oisd.nl/basic/",
  "https://adaway.org/hosts.txt",
  "https://pgl.yoyo.org/adservers/serverlist.php?hostformat=adblockplus&showintro=1&mimetype=plaintext",
  "https://raw.githubusercontent.com/nextdns/cname-cloaking-blocklist/master/domains",
}

var MediumAdditions = []string{
  "https://raw.githubusercontent.com/lightswitch05/hosts/master/docs/lists/facebook-extended.txt",
  "https://someonewhocares.org/hosts/zero/hosts",
  "https://urlhaus.abuse.ch/downloads/hostfile/",
}
```

```
49
50 var HeavyAdditions = []string{
51   "https://raw.githubusercontent.com/StevenBlack/hosts/master/hosts",
52   "https://raw.githubusercontent.com/Perflyst/PiHoleBlocklist/master/SmartTV-AGH.txt",
53   "https://raw.githubusercontent.com/DandelionSprout/adfilt/master/GameConsoleAdblockList.txt",
54 }
55
56 var IpBlockLists = []string{
57   "https://feodotracker.abuse.ch/downloads/ipblocklist.txt",
58   "https://www.spamhaus.org/drop/drop.txt",
59   "https://rules.emergingthreats.net/fwrules/emerging-Block-IPs.txt",
60   "https://cinsscore.com/list/ci-badguys.txt",
61 }
62
63 func getProfiles() map[string][]string {
64   profiles := make(map[string][]string)
65
66   light := append([]string(nil), CoreLists...)
67   light = append(light, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/light.txt")
68   profiles["light"] = light
69
70   medium := append([]string(nil), CoreLists...)
71   medium = append(medium, MediumAdditions...)
72   medium = append(medium, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/multi.txt")
73   profiles["medium"] = medium
74
75   heavy := append([]string(nil), CoreLists...)
76   heavy = append(heavy, MediumAdditions...)
77   heavy = append(heavy, HeavyAdditions...)
78   heavy = append(heavy, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/pro.txt")
79   profiles["heavy"] = heavy
80
81   return profiles
82 }
83
84 func loadEnvProfile() string {
85   profile := "medium"
86
87   if bytesData, err := os.ReadFile(".env"); err == nil {
88     scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
89     for scanner.Scan() {
90       line := strings.TrimSpace(scanner.Text())
91       if line != "" && !strings.HasPrefix(line, "#") {
92         parts := strings.SplitN(line, "=", 2)
93         if len(parts) == 2 {
94           key := strings.TrimSpace(parts[0])
95           val := strings.TrimSpace(parts[1])
96           if key == "GATEWAY_RULE_PROFILE" {
97             val = strings.Trim(val, "\"'")
98             profile = strings.ToLower(val)
99           }
100        }
101      }
102    }
103  } else if !os.IsNotExist(err) {
104    fmt.Printf("Warning: Failed to read .env file for profile configuration (%v). Using default.\n", err)
105  }
106
107  if envVal := os.Getenv("GATEWAY_RULE_PROFILE"); envVal != "" {
108    profile = strings.ToLower(envVal)
109  }
110
111  profiles := getProfiles()
112  if _, exists := profiles[profile]; !exists {
113    fmt.Printf("Warning: Profile '%s' is invalid. Falling back to 'medium'.\n", profile)
114    profile = "medium"
115  }
116  return profile
117 }
118
119 func loadDomains(filepath string) map[string]bool {
120   domains := make(map[string]bool)
121   if bytesData, err := os.ReadFile(filepath); err == nil {
122     scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
123     for scanner.Scan() {
124       line := strings.TrimSpace(scanner.Text())
125       if line != "" && !strings.HasPrefix(line, "#") {
126         domains[strings.ToLower(line)] = true
127       }
128     }
129   }
```

```go
  return domains
}

func parseDomainsFromContent(content string) map[string]bool {
  domains := make(map[string]bool)
  scanner := bufio.NewScanner(strings.NewReader(content))
  ipRegex := regexp.MustCompile(`^\d{1,3}\.\d{1,3}\.\d{1,3}\.\d{1,3}$`)
  agRegex := regexp.MustCompile(`^\|\|([a-z0-9.-]+)\^$`)
  domainRegex := regexp.MustCompile(`^[a-z0-9.-]+$`)

  for scanner.Scan() {
    line := strings.TrimSpace(scanner.Text())
    if line == "" || strings.HasPrefix(line, "!") || strings.HasPrefix(line, "[") {
      continue
    }
    parts := strings.SplitN(line, "#", 2)
    line = strings.TrimSpace(parts[0])
    if line == "" {
      continue
    }

    fields := strings.Fields(line)
    var rule string
    if len(fields) >= 2 {
      if ipRegex.MatchString(fields[0]) {
        rule = fields[1]
      } else {
        rule = fields[0]
      }
    } else {
      rule = fields[0]
    }

    rule = strings.TrimSpace(strings.ToLower(rule))

    var domain string
    if m := agRegex.FindStringSubmatch(rule); m != nil {
      domain = m[1]
    } else if domainRegex.MatchString(rule) {
      domain = rule
    } else {
      domain = rule
    }

    if domain != "" && domain != "localhost" && domain != "0.0.0.0" && domain != "127.0.0.1" && domain != "
     broadcasthost" {
      domains[domain] = true
    }
  }
  return domains
}

// WARNING: Parsing AdGuardHome.yaml using regex instead of a proper YAML parser
// is brittle and assumes the exact format currently emitted by AdGuard.
// If an AdGuard version bump updates the configuration schema format,
// this parser will fail silently or loudly and should be refactored to use gopkg.in/yaml.v3.
func getAdguardNativeLists() []string {
  yamlPath := "adguard/conf/AdGuardHome.yaml"
  var enabledUrls []string
  if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
    return enabledUrls
  }

  contentBytes, err := os.ReadFile(yamlPath)
  if err != nil {
    fmt.Printf("Warning: Failed to parse AdGuardHome.yaml for native lists (%v)\n", err)
    return enabledUrls
  }
  content := string(contentBytes)

  filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*?)(?:^[a-zA-Z_]+:|\z)`)
  match := filtersRegex.FindStringSubmatch(content)
  if len(match) > 1 {
    filtersText := match[1]
    blocks := strings.Split(filtersText, "\n  - ")
    urlRegex := regexp.MustCompile(`url:\s*(\S+)`)

    for _, block := range blocks {
      if strings.TrimSpace(block) == "" {
        continue
      }
```

```
      if strings.Contains(block, "enabled: true") {
        urlMatch := urlRegex.FindStringSubmatch(block)
        if len(urlMatch) > 1 {
          url := urlMatch[1]
          if !strings.Contains(url, "compiled_rules.txt") {
            enabledUrls = append(enabledUrls, url)
          }
        }
      }
    }
  }
  return enabledUrls
}

func getAdguardCompiledRulesCachePath() string {
  yamlPath := "adguard/conf/AdGuardHome.yaml"
  if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
    return ""
  }

  contentBytes, err := os.ReadFile(yamlPath)
  if err != nil {
    fmt.Printf("Warning: Failed to resolve compiled_rules.txt cache path (%v)\n", err)
    return ""
  }
  content := string(contentBytes)

  filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*?)(?:^[a-zA-Z_]+:|\z)`)
  match := filtersRegex.FindStringSubmatch(content)
  if len(match) > 1 {
    blocks := strings.Split(match[1], "\n  - ")
    idRegex := regexp.MustCompile(`id:\s*(\d+)`)
    for _, block := range blocks {
      if !strings.Contains(block, "compiled_rules.txt") {
        continue
      }
      idMatch := idRegex.FindStringSubmatch(block)
      if len(idMatch) > 1 {
        return fmt.Sprintf("adguard/work/data/filters/%s.txt", idMatch[1])
      }
    }
  }
  return ""
}

func handleFetchError(urlStr string, cacheFile string, err error) map[string]bool {
  if _, statErr := os.Stat(cacheFile); statErr == nil {
    fmt.Printf(" -> Warning: Failed to fetch %s (%v). Falling back to expired local cache.\n", urlStr, err)
    content, readErr := os.ReadFile(cacheFile)
    if readErr == nil {
      domains := parseDomainsFromContent(string(content))
      fmt.Printf(" -> Loaded from expired cache (%d domains).\n", len(domains))
      return domains
    }
    fmt.Printf(" -> Warning: Failed to read expired cache (%v). Using local lists only.\n", readErr)
  } else {
    fmt.Printf(" -> Warning: Failed to fetch %s (%v). Using local lists only for this source.\n", urlStr, err)
  }
  return make(map[string]bool)
}

func fetchRemoteDomains(urlStr string) map[string]bool {
  os.MkdirAll(CacheDir, 0755)

  hash := md5.Sum([]byte(urlStr))
  urlHash := hex.EncodeToString(hash[:])
  cacheFile := filepath.Join(CacheDir, fmt.Sprintf("%s.txt", urlHash))

  if stat, err := os.Stat(cacheFile); err == nil {
    fileAge := time.Since(stat.ModTime()).Seconds()
    if fileAge < 86400 {
      fmt.Printf("Loading remote blacklist from cache: %s\n", urlStr)
      content, err := os.ReadFile(cacheFile)
      if err == nil {
        domains := parseDomainsFromContent(string(content))
        fmt.Printf(" -> Loaded from cache (copy is %.1f hours old, %d domains).\n", fileAge/3600.0, len(domains
    ))
        return domains
      }
      fmt.Printf(" -> Warning: Failed to read cache file %s (%v). Re-fetching...\n", cacheFile, err)
    }
```

```
  }

  fmt.Printf("Fetching remote blacklist from: %s ...\n", urlStr)

  req, err := http.NewRequest("GET", urlStr, nil)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }
  req.Header.Set("User-Agent", "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7)")

  client := &http.Client{Timeout: 15 * time.Second}
  resp, err := client.Do(req)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }
  defer resp.Body.Close()

  if resp.StatusCode != 200 {
    return handleFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
  }

  bodyBytes, err := io.ReadAll(resp.Body)
  if err != nil {
    return handleFetchError(urlStr, cacheFile, err)
  }

  content := string(bodyBytes)
  domains := parseDomainsFromContent(content)

  if len(domains) < 10 {
    return handleFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only found %d domains. Possible
      404 or bad URL", len(domains)))
  }

  err = os.WriteFile(cacheFile, bodyBytes, 0644)
  if err != nil {
    fmt.Printf(" -> Warning: Failed to save cache file (%v)\n", err)
  }

  fmt.Printf(" -> Successfully fetched and cached %d domains.\n", len(domains))
  return domains
}

func optimizeSubdomains(domains map[string]bool, listName string) map[string]bool {
  fmt.Printf("Optimizing %s by removing redundant subdomains...\n", listName)
  rawCount := len(domains)
  if rawCount == 0 {
    return make(map[string]bool)
  }

  optimized := make(map[string]bool)
  for domain := range domains {
    if strings.ContainsAny(domain, "|^@$/") {
      optimized[domain] = true
      continue
    }

    parts := strings.Split(domain, ".")
    hasParent := false

    for i := 1; i < len(parts)-1; i++ {
      parent := strings.Join(parts[i:], ".")
      if domains[parent] {
        hasParent = true
        break
      }
    }

    if !hasParent {
      optimized[domain] = true
    }
  }

  saved := rawCount - len(optimized)
  reduction := float64(0)
  if rawCount > 0 {
    reduction = (float64(saved) / float64(rawCount)) * 100
  }
  fmt.Printf(" -> Reduced %s from %d to %d domains (-%d / %.1f%% reduction).\n", listName, rawCount, len(
     optimized), saved, reduction)
  return optimized
```

```
}

func parseIpsFromContent(content string) map[string]bool {
  ips := make(map[string]bool)
  scanner := bufio.NewScanner(strings.NewReader(content))

  for scanner.Scan() {
    line := strings.TrimSpace(scanner.Text())
    if line == "" || strings.HasPrefix(line, "#") || strings.HasPrefix(line, ";") {
      continue
    }

    fields := strings.Fields(line)
    if len(fields) > 0 {
      token := strings.TrimRight(fields[0], ";,")
      if _, err := netip.ParsePrefix(token); err == nil {
        ips[token] = true
      } else if _, err := netip.ParseAddr(token); err == nil {
        ips[token] = true
      }
    }
  }
  return ips
}

func handleIpFetchError(urlStr string, cacheFile string, err error) map[string]bool {
  if _, statErr := os.Stat(cacheFile); statErr == nil {
    fmt.Printf(" -> Warning: fetch failed (%v). Falling back to stale cache.\n", err)
    content, readErr := os.ReadFile(cacheFile)
    if readErr == nil {
      ips := parseIpsFromContent(string(content))
      fmt.Printf(" -> %d IPs loaded from stale cache.\n", len(ips))
      return ips
    }
    fmt.Printf(" -> Warning: stale cache also failed (%v).\n", readErr)
  } else {
    fmt.Printf(" -> Warning: %s failed (%v). No cache available. Skipping.\n", urlStr, err)
  }
  return make(map[string]bool)
}

func fetchRemoteIps(urlStr string) map[string]bool {
  os.MkdirAll(IpCacheDir, 0755)

  hash := md5.Sum([]byte(urlStr))
  urlHash := hex.EncodeToString(hash[:])
  cacheFile := filepath.Join(IpCacheDir, fmt.Sprintf("%s.txt", urlHash))

  if stat, err := os.Stat(cacheFile); err == nil {
    fileAge := time.Since(stat.ModTime()).Seconds()
    if fileAge < 86400 {
      fmt.Printf("Loading IP feed from cache: %s\n", urlStr)
      content, err := os.ReadFile(cacheFile)
      if err == nil {
        ips := parseIpsFromContent(string(content))
        fmt.Printf(" -> %d IPs (cache is %.1fh old).\n", len(ips), fileAge/3600.0)
        return ips
      }
      fmt.Printf(" -> Warning: cache read failed (%v). Re-fetching...\n", err)
    }
  }

  fmt.Printf("Fetching IP feed: %s ...\n", urlStr)
  req, err := http.NewRequest("GET", urlStr, nil)
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }
  req.Header.Set("User-Agent", "Mozilla/5.0")

  client := &http.Client{Timeout: 15 * time.Second}
  resp, err := client.Do(req)
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }
  defer resp.Body.Close()

  if resp.StatusCode != 200 {
    return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
  }

  bodyBytes, err := io.ReadAll(resp.Body)
```

```
  if err != nil {
    return handleIpFetchError(urlStr, cacheFile, err)
  }

  content := string(bodyBytes)
  ips := parseIpsFromContent(content)
  if len(ips) < 1 {
    return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only %d IPs found. Possible
     404", len(ips)))
  }

  err = os.WriteFile(cacheFile, bodyBytes, 0644)
  if err != nil {
    fmt.Printf(" -> Warning: cache write failed (%v)\n", err)
  }

  fmt.Printf(" -> %d IPs fetched and cached.\n", len(ips))
  return ips
}

func compileIpBlocklist() int {
  fmt.Println("\n IP Blocklist Compilation ")

  allIps := make(map[string]bool)
  var mu sync.Mutex
  var wg sync.WaitGroup

  maxThreads := 8
  if len(IpBlockLists) < 8 {
    maxThreads = len(IpBlockLists)
  }
  semaphore := make(chan struct{}, maxThreads)

  for _, urlStr := range IpBlockLists {
    wg.Add(1)
    go func(u string) {
      defer wg.Done()
      semaphore <- struct{}{}
      defer func() { <-semaphore }()

      ips := fetchRemoteIps(u)
      mu.Lock()
      for ip := range ips {
        allIps[ip] = true
      }
      mu.Unlock()
    }(urlStr)
  }
  wg.Wait()

  fmt.Printf("Total unique entries before filtering: %d\n", len(allIps))

  reservedStr := []string{
    "10.0.0.0/8",
    "172.16.0.0/12",
    "192.168.0.0/16",
    "127.0.0.0/8",
    "169.254.0.0/16",
    "100.64.0.0/10",
    "0.0.0.0/8",
  }
  var reserved []netip.Prefix
  for _, s := range reservedStr {
    p, _ := netip.ParsePrefix(s)
    reserved = append(reserved, p)
  }

  cleanIps := make(map[string]bool)
  for entry := range allIps {
    var prefix netip.Prefix
    p, err := netip.ParsePrefix(entry)
    if err != nil {
      addr, err2 := netip.ParseAddr(entry)
      if err2 != nil {
        continue
      }
      prefix = netip.PrefixFrom(addr, addr.BitLen())
    } else {
      prefix = p
    }

```

```
    overlaps := false
    for _, r := range reserved {
      if r.Overlaps(prefix) {
        overlaps = true
        break
      }
    }
    if !overlaps {
      cleanIps[entry] = true
    }
  }

  removed := len(allIps) - len(cleanIps)
  if removed > 0 {
    fmt.Printf(" -> Removed %d private/reserved entries.\n", removed)
  }
  fmt.Printf(" -> Final IP blocklist: %d entries.\n", len(cleanIps))

  os.MkdirAll(filepath.Dir(IpOutputPath), 0755)

  f, err := os.Create(IpOutputPath)
  if err != nil {
    fmt.Printf("Error writing IP output file: %v\n", err)
    return len(cleanIps)
  }
  defer f.Close()

  f.WriteString("# nullexit Compiled IP Blocklist\n")
  f.WriteString("# Sources: Feodo Tracker, Spamhaus DROP/EDROP, Emerging Threats, CINS\n")
  f.WriteString(fmt.Sprintf("# Entries: %d\n\n", len(cleanIps)))

  f.WriteString("create block_malicious_new hash:net maxelem 200000 -exist\n")

  var sortedIps []string
  for ip := range cleanIps {
    sortedIps = append(sortedIps, ip)
  }
  sort.Strings(sortedIps)

  for _, ip := range sortedIps {
    f.WriteString(fmt.Sprintf("add block_malicious_new %s -exist\n", ip))
  }

  f.WriteString("create block_malicious hash:net maxelem 200000 -exist\n")
  f.WriteString("swap block_malicious block_malicious_new\n")
  f.WriteString("destroy block_malicious_new\n")

  fmt.Printf("IP blocklist written to %s\n", IpOutputPath)
  return len(cleanIps)
}

func main() {
  profile := loadEnvProfile()
  fmt.Printf("Active memory profile: '%s'\n", strings.ToUpper(profile))

  localBlackList := loadDomains(BlackListPath)
  whiteList := loadDomains(WhiteListPath)

  blackList := make(map[string]bool)
  for k := range localBlackList {
    blackList[k] = true
  }

  profiles := getProfiles()
  urlsToFetch := profiles[profile]

  fmt.Printf("\nStarting concurrent downloads for %d lists...\n", len(urlsToFetch))
  startTime := time.Now()

  var mu sync.Mutex
  var wg sync.WaitGroup

  maxThreads := 16
  if len(urlsToFetch) < 16 {
    maxThreads = len(urlsToFetch)
  }
  semaphore := make(chan struct{}, maxThreads)

  for _, urlStr := range urlsToFetch {
    wg.Add(1)
    go func(u string) {
```

```
611         defer wg.Done()
612         semaphore <- struct{}{}
613         defer func() { <-semaphore }()
614
615         domains := fetchRemoteDomains(u)
616         mu.Lock()
617         for d := range domains {
618           blackList[d] = true
619         }
620         mu.Unlock()
621       }(urlStr)
622     }
623     wg.Wait()
624
625     fmt.Printf("Finished concurrent downloads in %.2f seconds using %d threads.\n", time.Since(startTime).Seconds
        (), maxThreads)
626
627     adguardNativeUrls := getAdguardNativeLists()
628     var adguardNativeDomains map[string]bool
629     if len(adguardNativeUrls) > 0 {
630       fmt.Printf("\nFetching %d AdGuard native list(s) for deduplication...\n", len(adguardNativeUrls))
631       adguardNativeDomains = make(map[string]bool)
632
633       maxNativeThreads := 4
634       if len(adguardNativeUrls) < 4 {
635         maxNativeThreads = len(adguardNativeUrls)
636       }
637       nativeSemaphore := make(chan struct{}, maxNativeThreads)
638
639       for _, urlStr := range adguardNativeUrls {
640         wg.Add(1)
641         go func(u string) {
642           defer wg.Done()
643           nativeSemaphore <- struct{}{}
644           defer func() { <-nativeSemaphore }()
645
646           domains := fetchRemoteDomains(u)
647           mu.Lock()
648           for d := range domains {
649             adguardNativeDomains[d] = true
650           }
651           mu.Unlock()
652         }(urlStr)
653       }
654       wg.Wait()
655
656       if len(adguardNativeDomains) > 0 {
657         fmt.Println("Deduplicating compiled rules against AdGuard native lists...")
658         originalSize := len(blackList)
659         for d := range adguardNativeDomains {
660           delete(blackList, d)
661         }
662         fmt.Printf(" -> Removed %d redundant rules already covered by AdGuard.\n", originalSize-len(blackList))
663       }
664     }
665
666     var contradictions []string
667     for d := range whiteList {
668       if blackList[d] {
669         contradictions = append(contradictions, d)
670       }
671     }
672
673     if len(contradictions) > 0 {
674       fmt.Printf("Detected %d contradictions. Whitelist taking priority.\n", len(contradictions))
675       sort.Strings(contradictions)
676       for _, domain := range contradictions {
677         if localBlackList[domain] {
678           fmt.Printf(" -> Removed local blacklist domain: %s\n", domain)
679         }
680         delete(blackList, domain)
681       }
682     }
683
684     blackList = optimizeSubdomains(blackList, "blacklist")
685     whiteList = optimizeSubdomains(whiteList, "whitelist")
686
687     os.MkdirAll(filepath.Dir(OutputPath), 0755)
688     outFile, err := os.Create(OutputPath)
689     if err != nil {
690       fmt.Printf("Error creating output file: %v\n", err)
```

```
    return
  }
  defer outFile.Close()

  outFile.WriteString("! Custom Compiled Rules (Auto-Generated)\n")
  outFile.WriteString(fmt.Sprintf("! Memory Profile: %s\n", strings.ToUpper(profile)))
  outFile.WriteString(fmt.Sprintf("! Total Block Rules: %d\n", len(blackList)))
  outFile.WriteString(fmt.Sprintf("! Native AdGuard Rules: %d\n", len(adguardNativeDomains)))
  outFile.WriteString(fmt.Sprintf("! Total Whitelist Rules: %d\n\n", len(whiteList)))

  outFile.WriteString("! --- Blacklist Rules ---\n")

  var sortedBlacklist []string
  for d := range blackList {
    sortedBlacklist = append(sortedBlacklist, d)
  }
  sort.Strings(sortedBlacklist)

  for _, domain := range sortedBlacklist {
    if strings.Contains(domain, "$") {
      parts := strings.SplitN(domain, "$", 2)
      dom, mod := parts[0], parts[1]
      if strings.HasPrefix(dom, "||") {
        if strings.HasSuffix(dom, "^") {
          outFile.WriteString(fmt.Sprintf("%s$%s\n", dom, mod))
        } else {
          outFile.WriteString(fmt.Sprintf("%s^$%s\n", dom, mod))
        }
      } else {
        outFile.WriteString(fmt.Sprintf("||%s^$%s\n", dom, mod))
      }
    } else if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|"
     ) || strings.HasSuffix(domain, "^") {
      outFile.WriteString(fmt.Sprintf("%s\n", domain))
    } else {
      outFile.WriteString(fmt.Sprintf("||%s^\n", domain))
    }
  }

  outFile.WriteString("\n! --- Whitelist Rules ---\n")
  var sortedWhitelist []string
  for d := range whiteList {
    sortedWhitelist = append(sortedWhitelist, d)
  }
  sort.Strings(sortedWhitelist)

  for _, domain := range sortedWhitelist {
    if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|") ||
     strings.HasSuffix(domain, "^") || strings.HasPrefix(domain, "@@") {
      rule := domain
      if !strings.HasPrefix(domain, "@@") {
        rule = "@@" + domain
      }
      outFile.WriteString(fmt.Sprintf("%s\n", rule))
    } else {
      outFile.WriteString(fmt.Sprintf("@@||%s^\n", domain))
    }
  }

  fmt.Printf("\nSuccessfully compiled %d block rules and %d allow rules to %s\n", len(blackList), len(whiteList
     ), OutputPath)

  cachedFilterPath := getAdguardCompiledRulesCachePath()
  if cachedFilterPath != "" {
    input, err := os.ReadFile(OutputPath)
    if err == nil {
      err = os.WriteFile(cachedFilterPath, input, 0644)
      if err == nil {
        fmt.Printf("Updated AdGuard filter cache: %s\n", cachedFilterPath)
      } else {
        fmt.Printf("Warning: Could not update AdGuard filter cache %s (%v)\n", cachedFilterPath, err)
      }
    } else {
      fmt.Printf("Warning: Could not read OutputPath to copy to cache (%v)\n", err)
    }
  } else {
    fmt.Println("Warning: Could not determine compiled_rules.txt cache path from AdGuardHome.yaml; skipping
     cache update.")
  }

  adguardReachable := false
```

```
768
769   dockerComposeYamlExists := false
770   if _, err := os.Stat("docker-compose.yml"); err == nil {
771     dockerComposeYamlExists = true
772   }
773   dockerPath, err := exec.LookPath("docker")
774   if err == nil && dockerComposeYamlExists {
775     cmd := exec.Command(dockerPath, "compose", "ps", "--status", "running", "-q", "adguardhome")
776     out, err := cmd.Output()
777     if err == nil && strings.TrimSpace(string(out)) != "" {
778       fmt.Println("Force-recreating AdGuard Home container to drop persisted filter state...")
779       upCmd := exec.Command(dockerPath, "compose", "up", "-d", "--force-recreate", "--no-deps", "adguardhome")
780       err := upCmd.Run()
781       if err == nil {
782         fmt.Println("AdGuard Home force-recreated successfully.")
783         time.Sleep(8 * time.Second)
784         adguardReachable = true
785       } else {
786         fmt.Println("Failed to restart AdGuard Home. Is it running?")
787       }
788     }
789   }
790
791   if adguardReachable {
792     agPass := "nullexit"
793     if bytesData, err := os.ReadFile(".env"); err == nil {
794       scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
795       for scanner.Scan() {
796         line := scanner.Text()
797         if strings.HasPrefix(line, "ADGUARD_PASSWORD=") {
798           parts := strings.SplitN(line, "=", 2)
799           if len(parts) == 2 {
800             agPass = strings.Trim(strings.TrimSpace(parts[1]), "\"'")
801           }
802           break
803         }
804       }
805     }
806
807     creds := base64.StdEncoding.EncodeToString([]byte(fmt.Sprintf("admin:%s", agPass)))
808     req, err := http.NewRequest("POST", "http://127.0.0.1:3000/control/filtering/refresh", bytes.NewBuffer([]
       byte("{}")))
809     if err == nil {
810       req.Header.Set("Authorization", fmt.Sprintf("Basic %s", creds))
811       req.Header.Set("Content-Type", "application/json")
812       client := &http.Client{Timeout: 15 * time.Second}
813       resp, err := client.Do(req)
814       if err == nil {
815         fmt.Printf("Triggered AdGuard filter refresh via REST API (HTTP=%d).\n", resp.StatusCode)
816         resp.Body.Close()
817       } else {
818         fmt.Printf("Warning: AdGuard filter refresh API call failed (%v). New whitelist will not load until
       next refresh tick (~24h) or manual refresh.\n", err)
819       }
820     }
821   }
822
823   compileIpBlocklist()
824 }
```

43 scripts/setup-common.sh

```
1  #!/bin/bash
2  #  1. Docker
3  step "Checking Docker"
4
5  if ! command -v docker >> output.log 2>&1; then
6      echo ""
7      if [[ "$OSTYPE" == "darwin"* ]]; then
8          echo "  Docker is not installed."
9          # Note: Homebrew aggressively rolls forward and discourages version pinning.
10         # This setup is confirmed stable on Colima v0.10.3 and Docker v29.6.1.
11         read -rp "  Would you like to automatically install Colima & Docker via Homebrew? [y/N]: "
       INSTALL_DOCKER
12         if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
13             step "Installing Colima and Docker..."
```

```
            brew install colima docker docker-compose
            step "Starting Colima VM..."
            colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged
        else
            echo ""
            echo "  Please install Docker manually. Two options:"
            echo "  A) Colima (Recommended): brew install colima docker docker-compose && colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged"
            echo "  B) Docker Desktop: https://www.docker.com/products/docker-desktop/"
            die "Install Docker, then re-run this script."
        fi
    else
        echo "  Docker is not installed."
        # Note: The Docker convenience script always fetches the latest stable release.
        # This setup is confirmed stable on Docker Engine v29.6.1.
        read -rp "  Would you like to automatically install Docker? (Requires sudo) [y/N]: " INSTALL_DOCKER
        if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
            step "Installing Docker..."
            curl -fsSL https://get.docker.com | sh
            sudo usermod -aG docker "$USER"
            echo -e "\n\033[0;32m[OK] Docker installed successfully!\033[0m"
            echo -e "\033[0;33m[!] CRITICAL: You must log out of your SSH session and log back in for Docker permissions to apply.\033[0m"
            die "Please log out, log back in, and re-run setup.sh."
        else
            echo "  Please install Docker manually:"
            echo "       curl -fsSL https://get.docker.com | sh"
            echo "       sudo usermod -aG docker \$USER && newgrp docker"
            die "Install Docker, then re-run this script."
        fi
    fi
fi

if ! docker info >> output.log 2>&1; then
    echo ""
    if [[ "$OSTYPE" == "darwin"* ]]; then
        echo "  Docker is installed but not running."
        echo "  If using Colima:        colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged"
        echo "  If using Docker Desktop: open the app."
    else
        echo "  Docker is installed but not running."
        echo "  Run: sudo systemctl start docker"
    fi
    die "Start Docker, then re-run this script."
fi

if ! docker compose version >> output.log 2>&1; then
    die "Docker Compose v2 not found. Update Docker or install the compose plugin:\n  https://docs.docker.com/compose/install/"
fi

ok "Docker $(docker --version | grep -oE '[0-9.]+' | head -1) is running."

#  1.5 Python 3
step "Checking Python 3"

if ! command -v python3 >> output.log 2>&1; then
    echo ""
    if [[ "$OSTYPE" == "darwin"* ]]; then
        echo "  Python 3 is not installed."
        echo "  macOS includes it via Xcode Command Line Tools, or you can install it via Homebrew:"
        echo "       brew install python3"
        die "Install Python 3, then re-run this script."
    else
        echo "  Python 3 is not installed."
        echo "  Please install it using your system's package manager. For example:"
        echo "       Debian/Ubuntu:  sudo apt install python3"
        echo "       Fedora:         sudo dnf install python3"
        echo "       Arch:           sudo pacman -S python"
        die "Install Python 3, then re-run this script."
    fi
fi

# Note: This setup is confirmed stable on Python 3.9.6 (macOS default).
PYTHON_VER=$(python3 --version 2>&1 | head -n 1)
ok "$PYTHON_VER is installed."

#  2. Tailscale (host)
step "Checking Tailscale (host)"
# The host needs its own Tailscale installation  separate from the Docker
```

```
# container  so Toggle-Gateway can run 'tailscale down' before stopping
# containers and 'tailscale up' after starting them. Without this, the
# exit-node deadlock that setup.sh is designed to prevent can still occur.

# Resolve the tailscale binary: CLI (brew/package) or bundled macOS app
TAILSCALE_BIN=""
if command -v tailscale >> output.log 2>&1; then
    TAILSCALE_BIN="tailscale"
elif [[ -x "/Applications/Tailscale.app/Contents/MacOS/Tailscale" ]]; then
    TAILSCALE_BIN="/Applications/Tailscale.app/Contents/MacOS/Tailscale"
fi

RECOMMENDED_TS_VERSION="1.98.5"

if [[ -z "$TAILSCALE_BIN" ]]; then
    warn "Tailscale not found on host  installing..."

    if [[ "$OSTYPE" == "darwin"* ]]; then
        if command -v brew >> output.log 2>&1; then
            step "Installing Tailscale CLI formula via Homebrew (standalone, no .app GUI)..."
            # Note: We install Tailscale via Homebrew, targeting the stable version (recommended: 1.98.5)
            brew install tailscale -q
            TAILSCALE_BIN="tailscale"

            step "Starting tailscaled as a per-user LaunchAgent (auto-starts at login)..."
            if brew services start tailscale 2>&1 | grep -qE 'Successfully|started'; then
                ok "tailscaled LaunchAgent registered."
            else
                warn "Could not auto-start tailscaled. Run 'brew services start tailscale' manually."
                warn "For a system-wide LaunchDaemon that runs before login, run instead:"
                warn "    sudo brew services start tailscale"
            fi
            echo ""
        else
            die "Homebrew not found. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION
    }):\n  https://tailscale.com/download/mac\n  Then re-run this script."
        fi

    elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
        curl -fsSL https://tailscale.com/install.sh | sh
        sudo systemctl enable --now tailscaled 2>> output.log || true
        TAILSCALE_BIN="tailscale"

    else
        die "Unsupported OS. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION}):\n
    https://tailscale.com/download\n  Then re-run this script."
    fi

    TAILSCALE_VER=$(${TAILSCALE_BIN} version 2>> output.log | head -n 1)
    if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
        warn "Tailscale installed, but version ($TAILSCALE_VER) differs from recommended version (
    $RECOMMENDED_TS_VERSION)."
    else
        ok "Tailscale installed (version $TAILSCALE_VER)."
    fi
else
    TAILSCALE_VER=$(${TAILSCALE_BIN} version 2>> output.log | head -n 1)
    if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
        warn "Tailscale is already installed ($TAILSCALE_VER), but recommended version is
    $RECOMMENDED_TS_VERSION for host/container consistency."
    else
        ok "Tailscale is already installed at recommended version ($TAILSCALE_VER)."
    fi
fi

# Authenticate the host machine if it isn't already connected.
# This is an interactive step  it opens a browser login URL.
# We do it now, before containers start, so Toggle-Gateway can
# safely run 'tailscale down/up' from day one.
if ! ${TAILSCALE_BIN} status >> output.log 2>&1; then
    echo ""
    echo "  Your host machine needs to join your Tailscale network."
    echo "  A browser window or login URL will appear  authenticate, then"
    echo "  return here. The script will continue automatically."
    echo ""
    if [[ "$OSTYPE" == "linux-gnu"* ]]; then
        sudo ${TAILSCALE_BIN} up
    else
        ${TAILSCALE_BIN} up
    fi
    ok "Host machine connected to Tailscale."
```

```
else
    ok "Host machine is already connected to Tailscale."
fi

#  3. wgcf
step "Checking wgcf (Cloudflare WARP key tool)"

if ! command -v wgcf >> output.log 2>&1; then
    warn "wgcf not found  installing..."

    if [[ "$OSTYPE" == "darwin"* ]]; then
        command -v brew >> output.log 2>&1 || die "Homebrew not found.\nInstall wgcf manually: https://github.com/ViRb3/wgcf/releases"
        brew install wgcf -q

    elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
        ARCH=$(uname -m)
        case $ARCH in
            x86_64)  WGCF_ARCH="amd64" ;;
            aarch64) WGCF_ARCH="arm64" ;;
            armv7l)  WGCF_ARCH="armv7" ;;
            *) die "Unsupported architecture: $ARCH.\nInstall wgcf manually: https://github.com/ViRb3/wgcf/releases" ;;
        esac

        echo "  Fetching latest release..."
        WGCF_VERSION=$(curl -sf https://api.github.com/repos/ViRb3/wgcf/releases/latest \
            | grep '"tag_name"' | cut -d'"' -f4)
        [[ -z "$WGCF_VERSION" ]] && die "Could not fetch wgcf version. Check your internet connection."

        sudo curl -sfL \
            "https://github.com/ViRb3/wgcf/releases/download/${WGCF_VERSION}/wgcf_${WGCF_VERSION#v}_linux_${WGCF_ARCH}" \
            -o /usr/local/bin/wgcf
        sudo chmod +x /usr/local/bin/wgcf
    else
        die "Unsupported OS. Install wgcf manually: https://github.com/ViRb3/wgcf/releases"
    fi
fi

ok "wgcf is ready."

#  4. WARP profile
step "Generating Cloudflare WARP profile"

if [[ -f "wgcf-profile.conf" ]]; then
    warn "wgcf-profile.conf already exists  skipping key generation."
else
    wgcf register --accept-tos 2>> output.log || die "wgcf register failed. Check your internet connection."
    wgcf generate 2>> output.log           || die "wgcf generate failed."
    ok "WARP profile generated."
fi

# Parse keys (IPv4 address only  gluetun doesn't need the IPv6 one)
PRIVATE_KEY=$(awk '/PrivateKey/{print $3}' wgcf-profile.conf)
PUBLIC_KEY=$(awk '/PublicKey/{print $3}'   wgcf-profile.conf)
ADDRESSES=$(awk '/Address/{print $3}'       wgcf-profile.conf | grep -v ':' | head -1)

[[ -z "$PRIVATE_KEY" ]] && die "Could not parse PrivateKey from wgcf-profile.conf"
[[ -z "$PUBLIC_KEY" ]]  && die "Could not parse PublicKey from wgcf-profile.conf"
[[ -z "$ADDRESSES" ]]   && die "Could not parse IPv4 Address from wgcf-profile.conf"

ok "WARP keys extracted."

#  5. Tailscale auth key
step "Tailscale authentication"

TS_AUTHKEY=""
if [[ -f ".env" ]]; then
    EXISTING_TS_KEY=$(read_env_var TS_AUTHKEY)
    if [[ -n "$EXISTING_TS_KEY" ]]; then
        TS_AUTHKEY="$EXISTING_TS_KEY"
        warn "Found existing TS_AUTHKEY in .env. Skipping prompt."
    fi
fi

if [[ -z "$TS_AUTHKEY" ]]; then
    echo ""
    echo "  Generate a key at: https://login.tailscale.com/admin/settings/keys"
    echo "   'Generate auth key'  enable Reusable if you plan to re-run setup"
    echo ""
```

```
    read -rp "  Paste your Tailscale auth key: " TS_AUTHKEY
    echo ""
fi

[[ -z "$TS_AUTHKEY" ]] && die "Tailscale auth key is required."
[[ "$TS_AUTHKEY" != tskey-auth-* ]] && warn "Key doesn't look like a Tailscale auth key  proceeding anyway."
ok "Auth key accepted."

#  6. AdGuard Home admin password
step "AdGuard Home admin account"

# Hardcoded default credentials to prevent lockout
# Username: admin
# Password: nullexit
if [[ -f ".env" ]]; then
    EXISTING_AG_PASS=$(read_env_var ADGUARD_PASSWORD)
    if [[ -n "$EXISTING_AG_PASS" ]]; then
        ADGUARD_PASSWORD="$EXISTING_AG_PASS"
        ADGUARD_PWD_ESC="$EXISTING_AG_PASS"
        ok "Found existing ADGUARD_PASSWORD in .env."
    else
        ADGUARD_PASSWORD="nullexit"
        ADGUARD_PWD_ESC="nullexit"
        ok "Default password set to 'nullexit'."
    fi
else
    ADGUARD_PASSWORD="nullexit"
    ADGUARD_PWD_ESC="nullexit"
    ok "Default password set to 'nullexit'."
fi


#  7. Write .env
step "Writing .env"

if [[ -f ".env" ]]; then
    read -rp "  .env already exists. Overwrite? [y/N]: " OVERWRITE
    if [[ "$OVERWRITE" != "y" && "$OVERWRITE" != "Y" ]]; then
        warn "Keeping existing .env."
    else
        WRITE_ENV=1
    fi
else
    WRITE_ENV=1
fi

if [[ "${WRITE_ENV:-0}" == "1" ]]; then
    # Preserve existing configuration flags if .env already exists
    EXISTING_PROFILE="medium"
    EXISTING_MSS="1120"
    EXISTING_HIJACK="true"
    EXISTING_EXIT="true"
    EXISTING_THRESH="6"
    EXISTING_KILL="false"
    EXISTING_BLOCKED="kp il"
    EXISTING_HOST_MTU="1200"

    if [[ -f ".env" ]]; then
        [[ -n "$(read_env_var GATEWAY_RULE_PROFILE)" ]] && EXISTING_PROFILE=$(read_env_var GATEWAY_RULE_PROFILE
    )
        [[ -n "$(read_env_var GATEWAY_MSS)" ]] && EXISTING_MSS=$(read_env_var GATEWAY_MSS)
        [[ -n "$(read_env_var GATEWAY_HIJACK_HOST)" ]] && EXISTING_HIJACK=$(read_env_var GATEWAY_HIJACK_HOST)
        [[ -n "$(read_env_var GATEWAY_USE_EXIT_NODE)" ]] && EXISTING_EXIT=$(read_env_var GATEWAY_USE_EXIT_NODE)
        [[ -n "$(read_env_var WARP_FAIL_THRESHOLD)" ]] && EXISTING_THRESH=$(read_env_var WARP_FAIL_THRESHOLD)
        [[ -n "$(read_env_var KILL_SWITCH)" ]] && EXISTING_KILL=$(read_env_var KILL_SWITCH)
        [[ -n "$(read_env_var BLOCKED_COUNTRIES)" ]] && EXISTING_BLOCKED=$(read_env_var BLOCKED_COUNTRIES)
        [[ -n "$(read_env_var HOST_MTU)" ]] && EXISTING_HOST_MTU=$(read_env_var HOST_MTU)
    fi

    cat > .env <<EOF
WIREGUARD_PRIVATE_KEY=${PRIVATE_KEY}
WIREGUARD_PUBLIC_KEY=${PUBLIC_KEY}
WIREGUARD_ADDRESSES=${ADDRESSES}
TS_AUTHKEY=${TS_AUTHKEY}
GATEWAY_RULE_PROFILE=${EXISTING_PROFILE}
# Safe MSS for double-tunneled traffic (Tailscale + WARP). Change to 1180 if you experience slow speeds and
    have a healthy path.
GATEWAY_MSS=${EXISTING_MSS}
# Lower the host Tailscale MTU to 1200 to prevent fragmentation when passing through the 1280-byte WARP tunnel.
HOST_MTU=${EXISTING_HOST_MTU}
```

```
# Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own internet is NOT hijacked).
GATEWAY_HIJACK_HOST=${EXISTING_HIJACK}
GATEWAY_USE_EXIT_NODE=${EXISTING_EXIT}
WARP_FAIL_THRESHOLD=${EXISTING_THRESH}
KILL_SWITCH=${EXISTING_KILL}
ADGUARD_PASSWORD=${ADGUARD_PASSWORD}
BLOCKED_COUNTRIES="${EXISTING_BLOCKED}"
EOF
    ok ".env written."
fi

#  8. Prepare directories
step "Preparing AdGuard directories"

mkdir -p adguard/conf adguard/work/userfilters

# Remove empty AdGuardHome.yaml that would cause a crash loop (same logic as Toggle-Gateway scripts)
if [[ -f "adguard/conf/AdGuardHome.yaml" && ! -s "adguard/conf/AdGuardHome.yaml" ]]; then
    rm "adguard/conf/AdGuardHome.yaml"
    warn "Removed empty AdGuardHome.yaml to prevent crash loop."
fi

ok "Directories ready."

#  9. Compile DNS filter rules
step "Compiling DNS filter rules (this may take a minute)"


#  10. Start containers
step "Starting containers"
# Note: tailscale depends on adguardhome being healthy (port 5335 up),
# so Docker Compose will hold it back automatically until the wizard is done.
docker compose up -d
ok "Containers starting."

#  11. Wait for AdGuard Home setup wizard
step "Waiting for AdGuard Home setup wizard"
echo "  (up to 60 seconds...)"

MAX=60; ELAPSED=0
until docker exec routing-fix curl -sf http://127.0.0.1:3000 >> output.log 2>&1; do
    sleep 2; ELAPSED=$((ELAPSED + 2))
    [[ $ELAPSED -ge $MAX ]] && die "AdGuard Home didn't come up.\nDebug with: docker compose logs adguardhome"
done

ok "AdGuard Home is up."

#  12. Configure AdGuard Home via API
step "Configuring AdGuard Home"

# Run all API calls in a single docker exec sh -c so the session cookie file
# is shared across calls without leaving the container.
docker exec routing-fix sh -c "
  set -e

  # Step 1: Complete the setup wizard (sets admin creds + DNS port to 5335)
  curl -sf -X POST http://127.0.0.1:3000/control/install/configure \
    -H 'Content-Type: application/json' \
    -d '{\"web\":{\"ip\":\"0.0.0.0\",\"port\":3000},\"dns\":{\"ip\":\"0.0.0.0\",\"port\":5335},\"username\":\"admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
    >/dev/null || true

  # Step 2: Wait for AdGuard to restart after wizard
  sleep 4
  WAITED=0
  until curl -sf http://127.0.0.1:3000 >/dev/null 2>&1; do
    sleep 2; WAITED=\$((WAITED + 2))
    [ \$WAITED -ge 30 ] && { echo 'AdGuard did not restart in time'; exit 1; }
  done

  # Step 3: Login to get session cookie
  curl -sf -X POST http://127.0.0.1:3000/control/login \
    -H 'Content-Type: application/json' \
    -d '{\"name\":\"admin\",\"password\":\"${ADGUARD_PWD_ESC}\"}' \
    -c /tmp/agh.cookie >/dev/null

  # Step 4: Point upstream DNS back through the WARP tunnel
  curl -sf -X POST http://127.0.0.1:3000/control/dns_config \
    -H 'Content-Type: application/json' \
    -b /tmp/agh.cookie \
```

```
    -d '{\"upstream_dns\":[\"127.0.0.1:53\",\"[/onion/]127.0.0.1:5353\"],\"bootstrap_dns\":[\"1.1.1.1\",\"
     8.8.8.8\"],\"upstream_mode\":\"load_balance\"}' \
    >/dev/null

  # Step 5: Register the compiled rules file as a filter list
  curl -sf -X POST http://127.0.0.1:3000/control/filtering/add_url \
    -H 'Content-Type: application/json' \
    -b /tmp/agh.cookie \
    -d '{\"name\":\"Custom Compiled Rules\",\"url\":\"/opt/adguardhome/work/userfilters/compiled_rules.txt\"}'
     \
    >/dev/null || true

  # Step 6: Force a filter refresh to load the rules immediately
  curl -sf -X POST http://127.0.0.1:3000/control/filtering/refresh \
    -H 'Content-Type: application/json' \
    -b /tmp/agh.cookie \
    -d '{\"whitelist\":false}' \
    >/dev/null || true

  rm -f /tmp/agh.cookie
"

ok "AdGuard Home configured."

#  13. Wait for Tailscale
step "Waiting for Tailscale to authenticate"
echo "  (Tailscale only starts once AdGuard's healthcheck passes  up to 90s)"

MAX=90; ELAPSED=0; TS_IP=""
until [[ -n "$TS_IP" ]]; do
    TS_IP=$(docker exec tailscale tailscale ip -4 2>> output.log || true)
    if [[ -z "$TS_IP" ]]; then
        sleep 3; ELAPSED=$((ELAPSED + 3))
        if [[ $ELAPSED -ge $MAX ]]; then
            warn "Tailscale hasn't authenticated yet."
            warn "Once it does, re-run the toggle script  it will resolve the IP dynamically."
            break
        fi
    fi
done

if [[ -n "$TS_IP" ]]; then
    ok "Tailscale IP: ${TS_IP} (resolved dynamically by the toggle script on each startup)"
fi
```

44 scripts/unlock-files.sh

```
#!/bin/bash
# Resolve project root relative to this script's location, so the script
# works regardless of the caller's CWD.
SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"

echo "Unlocking .env..."
if [ -f "$PROJECT_ROOT/.env" ]; then
  sudo cat "$PROJECT_ROOT/.env" > "$PROJECT_ROOT/.env.tmp" && mv "$PROJECT_ROOT/.env.tmp" "$PROJECT_ROOT/.env"
  echo " .env unlocked"
fi

echo "Unlocking AdGuard cache files..."
for f in "$PROJECT_ROOT"/adguard/work/userfilters/compiled_rules.txt \
         "$PROJECT_ROOT"/adguard/work/userfilters/ip_blocklist.ipset \
         "$PROJECT_ROOT"/adguard/work/userfilters/cache/*.txt \
         "$PROJECT_ROOT"/adguard/work/userfilters/cache/ip/*.txt \
         "$PROJECT_ROOT"/adguard/work/data/filters/*.txt; do
  if [ -f "$f" ]; then
    cat "$f" > "$f.tmp" && mv "$f.tmp" "$f"
    echo " Unlocked $f"
  fi
done

echo ""
echo "All locked files have been successfully bypassed via atomic rename!"
```

45 tor-bridges.txt

46 white_list.txt

```
0.client-channel.google.com
1drv.com
2.android.pool.ntp.org
akamaihd.net
akamaitechnologies.com
akamaized.net
amazonaws.com
android.clients.google.com
api-adservices.apple.com
api.ipify.org
api.rlje.net
app-api.ted.com
appleid.apple.com
apps.skype.com
appsbackup-pa.clients6.google.com
appsbackup-pa.googleapis.com
appspot-preview.l.google.com
apt.sonarr.tv
aspnetcdn.com
attestation.xboxlive.com
ax.phobos.apple.com.edgesuite.net
books-analytics-events.apple.com
brightcove.net
c.s-microsoft.com
cdn.cloudflare.net
cdn.embedly.com
cdn.optimizely.com
cdn.vidible.tv
cdn2.optimizely.com
cdn3.optimizely.com
cdnjs.cloudflare.com
cert.mgt.xboxlive.com
clientconfig.passport.net
clients1.google.com
clients2.google.com
clients3.google.com
clients4.google.com
clients5.google.com
clients6.google.com
connectivitycheck.android.com
connectivitycheck.gstatic.com
continuum.dds.microsoft.com
cpms.spop10.ams.plex.bz
cpms35.spop10.ams.plex.bz
cse.google.com
ctldl.windowsupdate.com
cws.conviva.com
d2c8v52ll5s99u.cloudfront.net
d2gatte9o95jao.cloudfront.net
dashboard.plex.tv
dataplicity.com
def-vef.xboxlive.com
delivery.vidible.tv
dev.virtualearth.net
device.auth.xboxlive.com
display.ugc.bazaarvoice.com
displaycatalog.mp.microsoft.com
dl.delivery.mp.microsoft.com
dl.dropbox.com
dl.dropboxusercontent.com
dns.msftncsi.com
download.sonarr.tv
drift.com
driftt.com
dynupdate.no-ip.com
ecn.dev.virtualearth.net
edge.api.brightcove.com
eds.xboxlive.com
fonts.gstatic.com
forums.sonarr.tv
```

```
g.live.com
geo-prod.do.dsp.mp.microsoft.com
geo3.ggpht.com
gfwsl.geforce.com
giphy.com
github.com
github.io
googleapis.com
gravatar.com
gstatic.com
help.ui.xboxlive.com
hls.ted.com
i.ytimg.com
i1.ytimg.com
iadsdk.apple.com
imagesak.secureserver.net
img.vidible.tv
imgix.net
imgs.xkcd.com
instantmessaging-pa.googleapis.com
intercom.io
jquery.com
jsdelivr.net
keystone.mwbsys.com
lastfm-img2.akamaized.net
licensing.xboxlive.com
live.com
livepassdl.conviva.com
login.live.com
login.microsoftonline.com
manifest.googlevideo.com
meta-db-worker02.pop.ric.plex.bz
meta.plex.bz
meta.plex.tv
microsoftonline.com
msftncsi.com
my.plexapp.com
nexusrules.officeapps.live.com
nine.plugins.plexapp.com
no-ip.com
node.plexapp.com
notes-analytics-events.apple.com
notify.xboxlive.com
npr-news.streaming.adswizz.com
ns1.dropbox.com
ns2.dropbox.com
o1.email.plex.tv
o2.sg0.plex.tv
ocsp.apple.com
office.com
office.net
office365.com
officeclient.microsoft.com
om.cbsi.com
onedrive.live.com
outlook.live.com
outlook.office365.com
pings.conviva.com
placehold.it
placeholdit.imgix.net
players.brightcove.net
pricelist.skype.com
products.office.com
proxy.plex.bz
proxy.plex.tv
proxy02.pop.ord.plex.bz
pubsub.plex.bz
pubsub.plex.tv
raw.githubusercontent.com
redirector.googlevideo.com
res.cloudinary.com
s.gateway.messenger.live.com
s.marketwatch.com
s.youtube.com
s.ytimg.com
s1.wp.com
s2.youtube.com
s3.amazonaws.com
sa.symcb.com
secure.avangate.com
secure.brightcove.com
```

```
secure.surveymonkey.com
services.sonarr.tv
skyhook.sonarr.tv
spclient.wg.spotify.com
ssl.p.jwpcdn.com
staging.plex.tv
status.plex.tv
t.co
t0.ssl.ak.dynamic.tiles.virtualearth.net
t0.ssl.ak.tiles.virtualearth.net
tawk.to
tedcdn.com
themoviedb.com
thetvdb.com
tinyurl.com
title.auth.xboxlive.com
title.mgt.xboxlive.com
traffic.libsyn.com
tvdb2.plex.tv
tvthemes.plexapp.com
twimg.com
ui.skype.com
video-stats.l.google.com
videos.vidible.tv
vidtech.cbsinteractive.com
weather-analytics-events.apple.com
widget-cdn.rpxnow.com
win10.ipv6.microsoft.com
wp.com
ws.audioscrobbler.com
www.dataplicity.com
www.googleapis.com
www.msftconnecttest.com
www.msftncsi.com
www.no-ip.com
www.youtube-nocookie.com
xbox.ipv6.microsoft.com
xboxexperiencesprod.experimentation.xboxlive.com
xflight.xboxlive.com
xkms.xboxlive.com
xsts.auth.xboxlive.com
youtu.be
youtube-nocookie.com
yt3.ggpht.com
zee.cws.conviva.com

# YouTube family  base domains so every subdomain (video stream servers on
# *.googlevideo.com, image hosts on *.ytimg.com, profile thumbnails on
# *.ggpht.com, etc.) is automatically allowlisted via the AdGuard '||domain^'
# wildcard. The specific s.youtube.com / manifest.googlevideo.com entries
# above are kept for clarity but will be optimized away by rule-compiler
# because their parents are now in this list.
youtube.com
googlevideo.com
ytimg.com
ggpht.com
gvt1.com
gvt2.com
gvt3.com
gvt1.net
gvt2.net
gvt3.net

edge-mqtt.facebook.com
gateway.instagram.com
i.instagram.com

facebook.com
fb.com
fbcdn.net
fbsbx.com
messenger.com
whatsapp.com
web.whatsapp.com
instagram.com
cdninstagram.com
static.whatsapp.net
media-bos5-1.cdn.whatsapp.net
media-yyz1-1.cdn.whatsapp.net
webtp.whatsapp.net
mmg.whatsapp.net
```