

# Nullexit: Unified Project Document

Omar Alaaeldein

July 17, 2026

## Contents

|           |                                    |            |
|-----------|------------------------------------|------------|
| <b>1</b>  | <b>Project Directory Tree</b>      | <b>3</b>   |
| <b>2</b>  | <b>README.md</b>                   | <b>5</b>   |
| <b>3</b>  | <b>agent.md</b>                    | <b>11</b>  |
| <b>4</b>  | <b>devref.md</b>                   | <b>19</b>  |
| <b>5</b>  | <b>diagrams.md</b>                 | <b>67</b>  |
| <b>6</b>  | <b>toggle.sh</b>                   | <b>72</b>  |
| <b>7</b>  | <b>recover.sh</b>                  | <b>100</b> |
| <b>8</b>  | <b>setup.sh</b>                    | <b>109</b> |
| <b>9</b>  | <b>scripts/setup-linux.sh</b>      | <b>110</b> |
| <b>10</b> | <b>docker-compose.yml</b>          | <b>111</b> |
| <b>11</b> | <b>scripts/common.sh</b>           | <b>114</b> |
| <b>12</b> | <b>scripts/watcher.sh</b>          | <b>127</b> |
| <b>13</b> | <b>scripts/crypto.sh</b>           | <b>131</b> |
| <b>14</b> | <b>.aiignore</b>                   | <b>132</b> |
| <b>15</b> | <b>.claude/settings.local.json</b> | <b>132</b> |
| <b>16</b> | <b>.gitignore</b>                  | <b>133</b> |
| <b>17</b> | <b>Recover-Gateway.applescript</b> | <b>133</b> |
| <b>18</b> | <b>Sweep-Gateway.applescript</b>   | <b>134</b> |
| <b>19</b> | <b>Toggle-Gateway.applescript</b>  | <b>134</b> |
| <b>20</b> | <b>benchmarks/benchmark.sh</b>     | <b>134</b> |
| <b>21</b> | <b>benchmarks/test_load.py</b>     | <b>135</b> |
| <b>22</b> | <b>black_list.txt</b>              | <b>136</b> |

|    |                                          |     |
|----|------------------------------------------|-----|
| 23 | cloud-init.yaml                          | 137 |
| 24 | docker/tor/Dockerfile                    | 138 |
| 25 | docker/tor/entrypoint.sh                 | 138 |
| 26 | launchd/com.nullexit.wake-recovery.plist | 139 |
| 27 | post-rules.txt                           | 140 |
| 28 | scripts/Dockerfile.routing-fix           | 140 |
| 29 | scripts/Dockerfile.rule-compiler         | 141 |
| 30 | scripts/Toggle-Gateway.bat               | 141 |
| 31 | scripts/check__circuits.py               | 143 |
| 32 | scripts/diagnose-host-leak.sh            | 145 |
| 33 | scripts/dns-proxy.py                     | 154 |
| 34 | scripts/fix-docker-bridge-collision.sh   | 155 |
| 35 | scripts/fix-ssh-delay.sh                 | 160 |
| 36 | scripts/generate__tex.py                 | 161 |
| 37 | scripts/logger/go.mod                    | 163 |
| 38 | scripts/logger/main.go                   | 163 |
| 39 | scripts/pf.conf                          | 167 |
| 40 | scripts/reinstall-host-tailscale.sh      | 167 |
| 41 | scripts/routing-fix.sh                   | 177 |
| 42 | scripts/rule-compiler/go.mod             | 180 |
| 43 | scripts/rule-compiler/main.go            | 181 |
| 44 | scripts/setup-common.sh                  | 191 |
| 45 | scripts/sweep.sh                         | 197 |
| 46 | scripts/unlock-files.sh                  | 201 |
| 47 | tor-bridges.txt                          | 201 |
| 48 | white_list.txt                           | 202 |

# 1 Project Directory Tree

```
nullexit/
|-- .aiignore
|-- .claude
|   |-- settings.local.json
|-- .gitignore
|-- LICENSE
|-- README.md
|-- Recover-Gateway.applescript
|-- Sweep-Gateway.applescript
|-- Toggle-Gateway.applescript
|-- agent.md
|-- benchmarks
|   |-- benchmark.sh
|   |-- test_load.py
|-- black_list.txt
|-- cloud-init.yaml
|-- devref.md
|-- diagrams.md
|-- docker
|   |-- tor
|       |-- Dockerfile
|       |-- entrypoint.sh
|-- docker-compose.yml
|-- launchd
|   |-- com.nullexit.wake-recovery.plist
|-- post-rules.txt
|-- recover.sh
|-- scripts
|   |-- Dockerfile.routing-fix
|   |-- Dockerfile.rule-compiler
|   |-- Toggle-Gateway.bat
|   |-- check_circuits.py
|   |-- common.sh
|   |-- crypto.sh
|   |-- diagnose-host-leak.sh
|   |-- dns-proxy.py
|   |-- fix-docker-bridge-collision.sh
|   |-- fix-ssh-delay.sh
|   |-- generate_tex.py
|   |-- logger
|       |-- go.mod
|       |-- main.go
|   |-- pf.conf
|   |-- reinstall-host-tailscale.sh
|   |-- routing-fix.sh
|   |-- rule-compiler
|       |-- go.mod
|       |-- main.go
|   |-- setup-common.sh
|   |-- setup-linux.sh
|   |-- sweep.sh
|   |-- unlock-files.sh
```

```
|  `-- watcher.sh
|-- setup.sh
|-- toggle.sh
|-- tor-bridges.txt
`-- white_list.txt
```

## 2 README.md

```
1 # nullexit: Tailscale + Cloudflare WARP Docker Gateway
2
3 > **Last updated:** July 12, 2026 [Architecture & Flow Diagrams ](./diagrams.md) [Full Dev Reference ](./
  devref.md)
4
5 **nullexit** is a chained network gateway that routes all Tailscale exit-node traffic through a Cloudflare WARP
  VPN tunnel double-encrypting every packet, hiding your ISP metadata, and providing network-wide DNS ad-
  blocking (AdGuard Home) and kernel-level IP threat blocking (`ipset`/`iptables`) for every device on your
  mesh.
6
7 ---
8
9 ## 1. Prerequisites
10 - Docker and Docker Compose installed.
11 - A Tailscale account.
12 - **macOS:** Install the standalone Tailscale CLI via Homebrew (`brew install tailscale`) and register `
  tailscaled` as a system service (`brew services start tailscale`). No `.app` GUI install required.
13 - **OS & Python:** Developed and tested against **macOS 26.5.2**. The project has zero external Python
  dependencies (no `requirements.txt`) and explicitly targets the default Apple-provided **Python 3.9.6** (`/
  usr/bin/python3`). While the scripts could function on newer versions, `nullexit` explicitly avoids
  experimenting with or modifying the built-in system Python environment to prevent unexpected OS-level side
  effects.
14
15 ## 2. Installation & Setup
16
17 Run the interactive setup script for your OS. It downloads `wgcf`, generates fresh Cloudflare WARP WireGuard
  keys, prompts for a Tailscale Auth Key, and writes your `.env`.
18
19 ```bash
20 ./setup.sh # macOS
21 ./scripts/setup-linux.sh # Linux
22 ```
23
24 ### Reference `.env`
25 This is a quick reference of the common variables. For the **complete, authoritative list** of every `.env`
  variable (including all optional/advanced settings), see `devref.md` 7`.
26
27 ```env
28 WIREGUARD_PRIVATE_KEY=<Generated>
29 WIREGUARD_PUBLIC_KEY=<Generated>
30 WIREGUARD_ADDRESSES=<Generated>
31 TS_AUTHKEY=tskey-auth-...
32 GATEWAY_RULE_PROFILE=medium # light | medium | heavy
33 # Safe MSS for double-tunneled traffic (Tailscale + WARP). Change to 1180 if you experience slow speeds and
  have a healthy path.
34 GATEWAY_MSS=1120
35 # Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own
  internet is NOT hijacked).
36 GATEWAY_HIJACK_HOST=true
37 GATEWAY_USE_EXIT_NODE=true
38 WARP_FAIL_THRESHOLD=6 # consecutive warp=off polls before auto-shutdown (default 6 = 30s)
39 KILL_SWITCH=false # enforce strict PF lock that breaks SSH if VPN fails
40 # On campus/enterprise networks (WPA2-Enterprise / 802.1x), AP Client Isolation blocks direct
  LAN traffic between devices. Tailscale attempts a local P2P upgrade anyway, which causes
41 # macOS gvproxy to SNAT-mangle the reply's source port poisoning the phone's WireGuard
  endpoint and blackholing 100% of its traffic. Setting this to false drops those local
42 # hole-punch packets so the phone safely falls back to Tailscale DERP relays.
43 # watcher.sh auto-detects 802.1x and AP isolation on every Wi-Fi change and overrides
  this setting automatically you only need to set this manually if auto-detection fails.
44 TAILSCALE_ALLOW_LAN_P2P=false # auto-managed; true only on trusted hotspots/home networks
45 ADGUARD_USER=admin # Username for AdGuard Home (defaults to admin if not set)
46 ADGUARD_PASSWORD=nullexit # Shared password for AdGuard Home and Tor ControlPort (defaults to nullexit
  if not set)
47 BLOCKED_COUNTRIES="kp il" # dynamically block country IP ranges via ipdeny.com
48 TOR_EXCLUDE_EXIT_NODES="" # comma-separated country codes to exclude as Tor exit nodes (e.g. {us},{gb
  })
49 DEBUG_TRACE=false # true = mirror a full command-by-command trace of the toggle to output.log
  (see 9)
50
51 ```
52
53 ## 3. Deploy the Gateway
54
55 **macOS (Recommended):**
56
57 ```bash
58 ./toggle.sh
59 ```
60 Running it again stops the gateway and restores your network. Handles the full lifecycle: Colima VM, containers
  , DNS hijacking, Tailscale exit-node routing, rule compilation, and sleep prevention.
61
```

```

62 **Linux / Manual:**
63 ```bash
64 docker compose up -d
65 ```
66
67 ## 4. Post-Deployment
68
69 ### Approve the Exit Node
70 1. Go to [Tailscale Admin Machines](https://login.tailscale.com/admin/machines).
71 2. Find the `tailscale` device `...` **Edit route settings** enable as Exit Node.
72
73 ### Verify Traffic Flow
74 On any Tailscale client device, select this gateway as the exit node, then visit [api.ipify.org](https://api.ipify.org) or run:
75 ```bash
76 curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
77 # expect: warp=on
78 ```
79
80 ### AdGuard DNS Setup
81
82 To ensure client devices correctly route DNS through the AdGuard container (and do not bypass it via `tailscaled`'s internal resolver), you **must** configure your Tailscale Admin Console:
83 1. Go to the **DNS** tab in the [Tailscale Admin Console](https://login.tailscale.com/admin/dns).
84 2. Enable **MagicDNS**.
85 3. Under **Nameservers**, add a **Custom** nameserver and enter the gateway's Tailscale IPv4 address (run `tailscale ip -4` on the gateway machine to find it, e.g., `100.x.x.x`).
86 4. Click the **three dots (...)** next to the nameserver you just added and enable **"Use with exit node"**.
87 5. Delete any other Global Nameservers (like Google or Cloudflare).
88 6. Toggle ON **Override local DNS**.
89
90 Traffic will now be correctly forced into the AdGuard sinkhole, and mobile devices will be blocked from bypassing it via encrypted DNS (DoH/DoT).
91
92 > [!WARNING]
93 > **Disable "Private DNS" on mobile devices.** Android's Private DNS (DoT/DoH) bypasses port 53 interception entirely. Go to `Settings Connections More connection settings Private DNS Off`.
94
95 > [!WARNING]
96 > **Bandwidth doubling:** Streaming 1GB of video on a remote device consumes 1GB download + 1GB upload on the host's network. Be mindful on metered connections.
97
98 ### AdGuard Dashboard
99 Navigate to `http://<gateway-tailscale-ip>:3000` credentials: **admin` / `nullexit`**.
100
101 ### Access Any Mesh Device From Anywhere (SFTP/SSH)
102
103 Once the gateway exit node is online, **any device on your Tailscale mesh can reach any other mesh device, anywhere in the world** as long as both devices are on the mesh and the target device has an SSH server running. No port forwarding, no public IPs, no firewall holes. You can SSH into any device using its MagicDNS hostname (e.g. `ssh username@my-macbook`) or its assigned Tailscale IP (`100.x.x.x`).
104
105 > **Security Tip (Zero Local Attack Surface):** It is highly recommended to disable macOS's native "Remote Login" and "Screen Sharing" in System Settings to prevent exposing those ports (22, 5900) to physical Wi-Fi networks (which scanners can fingerprint). The gateway enforces `tailscale up --ssh=true` and the `pf` Kill-Switch automatically blocks local Wi-Fi access to these ports. Since both MagicDNS and Tailscale IPs route through the exact same encrypted tunnel, they are equally secure but **MagicDNS is recommended simply for ease of use.**
106
107 ### Extreme OPSEC: The Secret Tor-over-WARP Proxy
108 `nullexit` includes a completely isolated, headless Tor infrastructure proxy designed for terminals and hacking tools (e.g., `curl`, `sqlmap`, `nmap`).
109 To protect against local malware fingerprinting:
110 1. **Procedural Ports:** The proxy ports are randomized into the ephemeral range (10000-65000) on every single boot using a strict kernel socket collision-avoidance check (`netstat`).
111 2. **Honey-Port Tripwire:** The gateway spawns a fake trap port. If any malware attempts a sequential port-scan on `localhost` to find the proxy, it trips the Honey-Port, instantly triggering a macOS desktop notification and a critical log alert.
112
113 > [!NOTE]
114 > **Threat Model limitation:** Port randomization and the Honey-Port only defeat blind, generic port scanners. They do not protect against targeted local malware that runs `lsof -i` or reads the `/tmp/nullexit-ports.env` file directly. To actually prevent malicious local processes from reaching the proxy, you must run a host-level outbound firewall (like LuLu 4.3.0+ or Little Snitch) with **localhost/loopback filtering explicitly enabled** to gate access by process identity.
115
116 To use the Tor proxy, run `cat /tmp/nullexit-ports.env` to find your `$TOR SOCKS_PORT` for the current session, and point your hardened browser (LibreWolf/Tor Browser) or CLI tool to `127.0.0.1:$TOR SOCKS_PORT`.
117
118 > [!WARNING]
119 > **Preventing DNS Leaks (SOCKS5 vs. SOCKS5-Hostname):**

```

```

120 > When routing command-line tools through Tor, you must force the tool to delegate DNS resolution to the
    Tor proxy. Using plain SOCKS5 (e.g., `socks5://` or curl's `--socks5` flag) will resolve hostnames locally
    * using the host's DNS settings (AdGuard/WARP) before establishing the connection. While this DNS traffic
    is encrypted inside the WARP tunnel, the destination domain name is still leaked to AdGuard and Cloudflare,
    undermining your anonymity.
121 >
122 > Use these correct protocols and configurations:
123 > - curl: Use `--socks5-hostname` or the `socks5h://` scheme (e.g., `curl -x socks5h://127.0.0.1:
    $TOR SOCKS_PORT https://check.torproject.org`). Do not use `--socks5` or `socks5://`.
124 > - sqlmap: Pass the proxy URL using the `socks5h://` scheme to ensure remote resolution: `--proxy="
    socks5h://127.0.0.1:$TOR SOCKS_PORT"`.
125 > - nmmap: Nmap's built-in `--proxies` option resolves hostnames locally. To safely scan targets without
    DNS leaks, tunnel it using `proxychains-ng` configured with `proxy_dns` enabled (add `socks5 127.0.0.1 <
    PORT>` to `/etc/proxychains.conf` or a local config, then run `proxychains4 nmap <args>`).
126
127
128 ##### Transparent .onion Browsing & Dark Web Access
129 In addition to the randomized SOCKS5 proxy port, `nullexit` provides seamless, transparent `.onion` domain
    browsing for all devices on your Tailscale mesh:
130 - Zero-Config DNS: AdGuard Home automatically intercepts queries for `.onion` domains and resolves them to
    Tor's virtual mapping subnet.
131 - Tailscale Auto-Routing: The virtual subnet is mapped to the `198.18.0.0/15` benchmark range (RFC 2544).
    Since this range is not a private IP subnet, Tailscale's Exit Node mode automatically routes it to the
    gateway (bypassing local LAN exclusions).
132 - Transparent NAT Proxying: Kernel firewall rules in the gateway redirect all TCP traffic destined for
    `198.18.0.0/15` directly into Tor's transparent port (`9040`).
133 - Exit Node Exclusion: Supports the ability to exclude specific countries from being used as Tor exit nodes
    via the `TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`.
134
135 ##### Enable in Web Browsers
136 To prevent accidental DNS leaks, web browsers block `.onion` lookups by default. To browse dark web sites
    natively:
137 - Firefox: Go to `about:config`, search for `network.dns.blockDotOnion`, and set it to `false`. Type any
    onion address (e.g. `duckduckgogg42xjoc72x3sjasowoarfbgcmvfimafitt6twagswzczad.onion`) directly into the URL
    bar and it will load.
138 - Mobile Browsers (iOS/Android): Connect to your Tailscale exit node, open Firefox (with blockDotOnion
    disabled), and browse onion sites natively on the go.
139
140 ##### Hardening: Using obfs4 Tor Bridges
141 If you want to disguise your Tor traffic from the WARP exit node provider, you can optionally enable Tor
    bridges:
142 1. Obtain private `obfs4` bridge lines from [bridges.torproject.org](https://bridges.torproject.org) or the
    official Telegram bot.
143 2. Create a file named `tor-bridges.txt` in the root of the `nullexit` folder and paste your bridge lines
    inside it (one per line).
144 3. Set `TOR_USE_BRIDGES=true` in your `.env` file and restart the gateway. The proxy will refuse to start if
    the bridges file is empty or missing.
145
146 ##### Example: Browse your Android phone's files from your Mac
147
148
149 1. On your Android phone, install [Termux](https://termux.dev/) and [Termux:Boot](https://f-droid.org/
    packages/com.termux.boot/) from F-Droid.
150 2. In Termux, install and start the SSH server:
151 ```bash
152 pkg install openssh
153 sshd
154 ```
155 Termux defaults to port 8022 (ports below 1024 require root on Android).
156 3. Find your phone's Tailscale IP:
157 ```bash
158 tailscale ip -4 # e.g. 100.87.42.87
159 ```
160 4. Stop Android from killing Termux do ALL of these (Samsung One UI is especially aggressive):
161 - Settings Apps Termux Battery Unrestricted
162 - Device Care Battery Background usage limits Sleeping apps / Deep sleeping apps remove Termux if
    listed
163 - In Termux, run `termux-wake-lock` to hold a persistent wakelock
164 5. On your Mac/PC, open [Cyberduck](https://cyberduck.io/) (or any SFTP client: WinSCP, FileZilla, FE File
    Explorer on iOS) and connect to:
165 - Server: `100.87.42.87` (your phone's Tailscale IP)
166 - Port: `8022`
167 - Protocol: SFTP (SSH File Transfer Protocol)
168 - Username: (your Termux username, usually `u0_a123` run `whoami` in Termux to check)
169 - Password: (your Termux password set with `passwd` in Termux if you haven't already)
170
171 That's it you can now browse, download, and upload files to your phone from any device on the mesh, no matter
    where either device is physically located. This works identically for any mesh device: Mac Mac, Windows
    Linux, phone NAS, etc.
172

```

```

173 > **Troubleshooting:** If the connection hangs for ~30 seconds before failing, your phone's SSH server isn't
      running. Open Termux and run `sshd` again. If Android keeps killing it despite Unrestricted battery, the
      termux-wake-lock` command is your fix.
174
175 ---
176
177 ## 5. Multi-Layer Firewall & Kill-Switch
178
179 ### Layer 1 DNS Sinkhole (AdGuard Home)
180 Blocks ads and tracking domains before they are downloaded. In automated Lighthouse tests on ad-heavy sites:
181 - **Time to Interactive:** ~22% faster (51.0s 41.8s)
182 - **Speed Index:** ~20% faster (15.3s 12.8s)
183
184 Customize blocking via `black_list.txt` and `white_list.txt`. Rule profiles (`light` / `medium` / `heavy`) are
      set in `.env`.
185
186 ### Layer 2 Kernel IP Firewall (`ipset`/`iptables`)
187 On every startup, the `rule-compiler` fetches threat-intelligence feeds (Spamhaus DROP, Feodo Tracker, Emerging
      Threats, CINS) and compiles **~16,700 unique malicious IPs/CIDRs** into the kernel `FORWARD` chain
      blocking both outbound C2 connections and inbound attack traffic. Zero configuration required.
188
189 ### Layer 3 Geo-IP Blocking
190 Dynamically blocks all traffic to and from specific countries using live IP ranges from `ipdeny.com`.
191 To add countries to your blocklist, open your `.env` file and set the `BLOCKED_COUNTRIES` variable with 2-
      letter ISO country codes (e.g. `BLOCKED_COUNTRIES="kp il cn ru"`), then restart the gateway.
192
193 ### Layer 4 macOS PF Kill-Switch
194 A native macOS Packet Filter (`pf`) kill-switch that enforces a strict default-deny policy on your physical Wi-
      Fi interface (`en0`). When `KILL_SWITCH=true` is set in your `.env`, it drops all outgoing traffic except
      the Cloudflare WARP endpoints and Tailscale DERP relays.
195
196 - **Failsafe Design:** If the VPN tunnel crashes or Docker dies, your Mac completely loses internet access
      rather than leaking your real IP.
197 - **Remote-Access Safe:** Carefully designed to whitelist Tailscale's control plane (`192.200.0.0/24`), `utun*`
      tunnel traffic, and local LAN subnets. This guarantees that your remote SSH sessions and local AirDrop
      will survive even when the kill switch engages.
198 - **Setup Requirement:** The toggle scripts need permission to manipulate the network and firewall in the
      background without prompting you for a password. You must configure your passwordless sudo config by
      running exactly:
199 ```bash
200 echo "$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /usr/bin/dscacheutil, /usr/bin/killall,
      /usr/bin/pkill, /bin/kill, /sbin/route, /sbin/ifconfig, /usr/bin/true, /opt/homebrew/bin/brew, /usr/local/
      bin/brew, /usr/bin/python3 -I $PWD/scripts/dns-proxy.py, /opt/homebrew/bin/python3 -I $PWD/scripts/dns-
      proxy.py, /usr/local/bin/python3 -I $PWD/scripts/dns-proxy.py" | sudo tee /etc/sudoers.d/nullexit
201 ```
202
203 ---
204
205 ## 6. Toggle Scripts
206
207 ### macOS `toggle.sh`
208 Fully automated lifecycle script. Also supports a native `.app` launcher:
209 ```bash
210 osacompile -o "Toggle Gateway.app" Toggle-Gateway.applescript
211 ```
212
213 Optional `.env` settings (common subset; the full canonical table lives in `devref.md`):
214 - `GATEWAY_BYPASS_PING=true` proceed even if pre-flight connectivity checks fail.
215 - `GATEWAY_USE_EXIT_NODE=false` DNS-only mode, skip exit-node routing.
216 - `GATEWAY_HIJACK_HOST=false` skip DNS hijacking on the host (provides VPN/adblocking only to external
      Tailscale peers).
217 - `GATEWAY_MSS=1180` override the TCP MSS clamp for the double-tunnel (default `1120`; raise to `1180` for
      more speed on a healthy path).
218 - `HOST_MTU=1200` override the host Tailscale interface MTU. Defaults to 1200 to fit inside the 1280-byte WARP
      tunnel.
219 - `KILL_SWITCH=true` enforce strict network lock that breaks SSH if the VPN fails.
220 - `STOP_COLIMA_ON_EXIT=true` fully shut down the Colima VM on toggle-off (saves battery on dedicated hosts).
221 - `TAILSCALE_ALLOW_LAN_P2P=true` allow direct Tailscale LAN P2P connections between devices on the same
      network.
222 - **Default: `false`** On campus or enterprise Wi-Fi (WPA2-Enterprise / 802.1x), AP Client Isolation blocks
      local device-to-device traffic. When Tailscale still attempts a local P2P upgrade, macOS's `gvproxy` layer
      SNAT-mangles the reply's source port. WireGuard's endpoint roaming treats this as a legitimate address
      change and updates the phone's outbound path to the newly-mangled port which has no listener. This
      blackholes 100% of the phone's traffic while the DERP relay path is abandoned. Setting this to `false`
      preemptively drops all RFC1918-destined Tailscale UDP from the gateway so the phone receives a clean
      failure and safely falls back to DERP.
223 - **Auto-managed:** `watcher.sh` automatically detects WPA2-Enterprise via `system_profiler SPAirPortDataType`
      (the `airport` binary was removed in macOS 14.4) and probes for AP isolation by pinging the default
      gateway (`route -n get 1.1.1.1`) on every network change. It writes the detected value to
      `lan_p2p_detected` in the repo root, which `routing-fix.sh` re-reads every 30 seconds **no restart required
** Only set this manually if auto-detection fails. See `devref.md 15.7.1` for the full SNAT endpoint

```



```

    poisoning analysis and `devref.md 15.11.6` for the `airport` removal background. *(Note: Direct P2P between
    the gateway container and the host Mac itself is always permitted via the Docker bridge to bypass DERP
    relays, regardless of this setting).*
224 - Set to `true` on trusted home networks or Windows/mobile hotspots (no AP isolation) for a faster, lower-
    latency direct path.
225 - `WARP_FAIL_THRESHOLD=3` number of consecutive failed checks before forcing a gateway teardown (default: 6
    checks = 30s).
226 - `WARP_ENDPOINT_1` / `WARP_ENDPOINT_2` override the Cloudflare WARP WireGuard endpoints.
227
228 ### Linux `toggle.sh` / `recover.sh`
229 ```bash
230 ./toggle.sh # toggle ON/OFF (cross-platform, dispatches on $OSTYPE)
231 ./recover.sh # recovery tool (cross-platform)
232 ```
233
234 ### Sleep Prevention (macOS)
235 When the gateway is ON, `caffeinate -i` runs in the background to prevent idle system sleep, keeping Docker and
    all services alive even on battery. The display can still sleep normally.
236
237 ### Auto-Recovery (macOS post-wake / post-roam)
238 Install the launchd LaunchAgent so the gateway self-heals after sleep/wake and Wi-Fi roams:
239 ```bash
240 cp ./launchd/com.nulllexit.wake-recovery.plist ~/Library/LaunchAgents/
241 sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" ~/Library/LaunchAgents/com.nulllexit.wake-recovery.plist
242 launchctl load -w ~/Library/LaunchAgents/com.nulllexit.wake-recovery.plist
243 ```
244 See `devref.md 15.5.1` for the full deep dive.
245
246 ---
247
248 ## 7. Networking Notes
249
250 - **Port conflicts:** None. Tailscale uses dynamic UDP ports; WARP uses outbound UDP:2408. The gateway binds
    `0.0.0.0:41642/udp` on the host for Tailscale's WireGuard listener this is intentional and required to
    receive inbound hole-punch packets from peers on the internet. All other internal ports (AdGuard :5335,
    SOCKS5 :1080) are bound to `127.0.0.1` only and are not reachable from the LAN.
251 - **Tailscale P2P vs DERP:** On the same Wi-Fi, Tailscale establishes a fast P2P connection. On cellular (CGNAT
    ), it always falls back to a DERP relay (~80-200ms). See `devref.md 5.1`.
252 - **AirDrop / AirPlay:** Unaffected the gateway only touches standard Wi-Fi/Ethernet interfaces, not `awdl0`.
253 - **VPN coexistence:** Do not run a local VPN client (WARP, Mullvad, NordVPN) simultaneously with the exit
    node. Both fight for the default route and will blackhole your connection.
254 - **Banking & financial sites:** May intermittently block logins through WARP. Banks use anti-fraud databases
    that flag datacenter IP ranges (like Cloudflare's `104.28.x.x`) as VPN/proxy traffic. Success fluctuates
    depending on which WARP IP you land on. If a site blocks you, either disable the exit node temporarily for
    that session or use the bank's mobile app (which often uses different fraud detection).
255 - **MSS clamping (Bufferbloat):** Double-tunnelling adds ~120-160 bytes of overhead per packet. This is now
    automatically handled natively via the macOS `pf` firewall, which universally applies TCP MSS Clamping to
    all outbound packets to strictly prevent MTU fragmentation and bufferbloat. No manual `.env` configuration
    is required. Additionally, the host's Tailscale interface MTU is explicitly lowered to 1200 (configurable
    via `HOST_MTU` in `.env`) to guarantee packets fit inside the container's WARP tunnel without fragmenting.
256
257 ---
258
259 ## 8. Privacy & Security Hardening
260
261 The stack shifts trust across four layers: WARP hides traffic from your ISP, AdGuard blocks tracking at DNS,
    Tailscale encrypts device-to-device, and `ipset` blocks known malicious IPs at the kernel.
262
263 For higher security, the gateway supports these drop-in upgrades:
264
265 | Layer | Default | Hardened |
266 |---|---|---|
267 | VPN Exit | Cloudflare WARP (US) | Mullvad (Swedish, no-logs, anonymous payment) |
268 | DNS Resolver | AdGuard via WARP | AdGuard via Mullvad DoH |
269 | Coordination | Tailscale (US) | Headscale (self-hosted) |
270 | Auth | Persistent OAuth keys | Ephemeral auth keys |
271 | Content | WireGuard asymmetric | WireGuard + Pre-Shared Keys (PSK) |
272
273 ##### Python Runtime Airgap (Isolated Mode)
274 **Never** install third-party `pip` packages (like `requests` or `pyyaml`) into the Python environment that
    `nullexit` uses. All python scripts in this project (e.g., the `dns-proxy.py` daemon) are strictly
    engineered to run on the zero-dependency Python Standard Library. To strictly isolate the environment from
    supply-chain attacks, `nullexit` executes these scripts using **Python's Isolated Mode** (`python3 -I`).
    This deliberately blinds the execution environment to any malicious user-installed `pip` packages on the
    host, permanently air-gapping the gateway from PyPI.
275
276 See `devref.md 9` for the full threat model, Snowden programme analysis, and trust assumptions.
277
278 ---
279
280 ## 9. Debugging

```

```

281 - **`output.log`** All stderr from `toggle.sh`, `recover.sh`, `setup.sh`, and the rule-compiler is written
282 here. The WARP Watcher (`WARP_FAIL_THRESHOLD`) also writes its events here `WARP DOWN`, `WARP RECOVERED`,
and `WARP SHUTDOWN` notices. Check this first.
283 - Every run now records lifecycle breadcrumbs regardless of `DEBUG_TRACE`: an `EXEC:` / `EXIT <code>:` line
brackets each heavyweight command (`colima start`, `docker compose up -d --build`, `docker compose down`),
and an `EXIT: code=<N> phase=<init|stop|start>` line is written when the script exits **for any reason**
including a killed terminal, `kill`, or OOM. This closes a prior blind spot where an interrupted **START**
(e.g. a `--restart` race while Wi-Fi was re-associating) left the log ending abruptly after the teardown
with no indication of what failed. If a toggle "does nothing" or the gateway ends up half-up, `grep -E '
EXEC:|EXIT ' output.log | tail` shows the last command run and its result.
284 - **`DEBUG_TRACE=true`** (in `.env`) Full command-by-command execution trace of `toggle.sh`, mirrored to `
output.log` for deep post-mortems of intermittent failures. Each traced line is timestamped and tagged with
`file:line:function` (via a custom `PS4`), so you can follow the *exact* code path a run took including
subshells, Docker/Colima calls, and where it aborted. Off by default because it is verbose (the log grows
quickly); enable it only while reproducing a specific problem, then set it back to `false`. Implementation
note: it prefers bash's `BASH_XTRACEFD` (bash 4.1, e.g. most Linux) to send the trace to the log while
keeping your terminal clean; on macOS's stock `/bin/bash` 3.2 it transparently falls back to tee'ing stderr
to the log (trace also appears in the terminal). It never uses the dynamic-fd (`exec {var}>>`) form, so it
is safe on 3.2. To read a trace: `grep -nE '\n ' output.log`.
285 - **`bash scripts/diagnose-host-leak.sh`** One-shot host-routing diagnostic. Runs 8 checks and classifies your
state into one of four known scenarios (System Extension conflict, SOCKS5 fallback, IPv6 leak, route-table
freeze) or OK, and prints the exact fix command. Check 5 uses `route -n get 1.1.1.1` to ask the kernel
which interface real traffic actually resolves through (more accurate than `netstat -rn | grep default` on
macOS). Pass `--fix` to apply the remediation automatically. Pass `--watch` (or `--watch 30`) to run a full
baseline diagnostic then continuously monitor warp/IPv6/default-route for leaks, alerting on any state
change.
286 - **`bash scripts/sweep.sh`** One-shot, read-only gateway health sweep. Runs 6 checks containers up/healthy,
double-tunnel/IP-leak (host egress == container egress with `warp-on`), direct P2P to the gateway (no DERP
relay), AdGuard compiled-rule counts + live DNS interception, PF kill-switch armed, and tailnet
reachability (a `tailscale ping` to every online peer) then rolls them into a PASS/WARN/FAIL verdict.
Writes a timestamped `sweep-<UTC>.txt` report to the repo root (diffable across runs) and exits non-zero on
FAIL, so it can gate launchd hooks or CI. Pass `--quiet` for just the verdict line. It never mutates
routing, PF, DNS, or containers safe to run against a live gateway at any time.
287 - **`bash scripts/fix-docker-bridge-collision.sh`** One-shot fix for Docker bridge IP conflicts. If your local
Wi-Fi router assigns you an IP in the `172.17.0.0/16` range, it will violently collide with Docker's
default internal bridge, permanently killing your internet. This script detects the collision and auto-
patches your Colima VM to use a safe subnet (e.g. `10.200.0.1/24`).
288 - **`devref.md`** Complete architecture reference, routing deep-dives, and full resolved-issues log. Read this
before modifying any routing or firewall logic.
289
290 > [!CAUTION]
291 > **Two infinite routing loops are possible.** (1) If a hotspot device providing internet to the gateway host
is *also* using the gateway as its exit node, packets bounce forever between the two. (2) When the host's
default route is redirected to `utun`, WARP endpoint packets can loop back into the tunnel. Both are
documented in `devref.md 15.2.1` and `15.2.2` and are automatically handled by `toggle.sh`.
292
293 ---
294
295 ## 10. Unified Project Document (Quine)
296
297 The project includes a `generate_tex.py` script that compiles the entire source code, configurations, and
documentation into a single, unified LaTeX document (`nullexit_unified.tex` / `.pdf`).
298
299 Funnily enough, this document acts as a "project-level quine." It encapsulates the entirety of its own source
code, its documentation, and the exact Python code required to generate itself. It serves as a self-
contained, long-term archiving strategy providing everything needed to reconstruct and understand the
`nullexit` system from a single file.
300
301 ---
302
303 ## 11. Acknowledgements
304 - **[SyameimaruKoa](https://github.com/SyameimaruKoa):** Dual-stack TCP MSS clamping, `SIGHUP` state-tracking
in the routing sidecar, and `TS_AUTH_ONCE` integration.
305
306 ## 12. License
307 GNU Affero General Public License v3. See [LICENSE](./LICENSE).
308
309 ## 13. Changelog
310
311 - **July 17, 2026** `scripts/sweep.sh` (the read-only automated health sweep added July 15, mirroring `sweep.
md`) gained a 6th check: **tailnet reachability** every Tailscale peer reported online is probed with
`tailscale ping` (TSMP, so it works even on phones that drop ICMP) and rolled into the PASS/WARN/FAIL
verdict. Documented in 9 above.
312 - **July 12, 2026** Tor ControlPort dynamic password authentication (unified with `ADGUARD_PASSWORD`),
enabling live circuit inspection via `/sweep`. See `devref.md 15.10.2`.
313 - **July 11, 2026** Tor container hardening: `obfs4proxy` restored (base image `debian:bookworm-slim`),
robust bridge-file validation, and image pinning. See `devref.md 15.10.1`.
314 - **July 10, 2026** Roaming/startup stabilization suite: fixed the Wi-Fi-roam host-IP leak (raw ISP `132.x`)
caused by the WARP Watcher racing post-wake recovery, plus a batch of pre-flight/concurrency/kill-switch/
Tailscale-flag bugs and the macOS SNAT endpoint-poisoning P2P blackhole. Full post-mortems in the Incident

```

```

Log see `devref.md 15.2.7`, `15.5.3`~`15.5.5`, `15.2.8`~`15.2.9`, `15.7.1`, `15.7.3`~`15.7.5`, `15.12.18`.
315 - **July 6, 2026** Edge-case security and stability hardening: DNS Watcher interface independence, robust lock
    -file PID verification, fail-closed crypto integrity check, strict `iptables` pinning for tailscale0tun0,
    Docker port binding to `127.0.0.1` to prevent LAN exposure, routing verification via `netstat -nr`. See `
    devref.md 17`.
316 - **July 4, 2026** Colima bridged networking, SOCKS5 proxy migrated natively to Tailscale, headless prompt
    crashes fixed, proxy bypass domains added to fix LAN P2P. Python dependencies completely eliminated: `
    logger` and `rule-compiler` ported to blazing-fast Go multi-stage Docker builds. Added `crypto.sh` to
    enforce cryptographic signature validation on all core bash scripts. Massively refactored `toggle.sh` and `
    recover.sh` to eliminate ~200 lines of duplicated code by migrating network and process logic to `scripts/
    common.sh`. Added a concurrency mutex to `toggle.sh`, resolved unbound `TS_AUTHKEY` check crashes on setup,
    and integrated Go validation to the CI pipeline.
317 - **July 3, 2026** **Critical Security Updates:** Addressed the overnight silent IP leak and improved geo-
    blocking. Introduced `scripts/host-leak-probe.sh` (a sub-second host-egress leak detector) and a
    statistical threshold auto-shutdown watcher. Added `unlock-files.sh` to safely reset file permissions.
    Resolved numerous startup regressions and system bugs.
318 - **July 2, 2026** Two infinite routing loops resolved (`devref.md 15.2.2`): host-VM tunnel loop (WARP bypass
    routes now via gateway IP + manual `utun*` redirection) and Docker Compose subnet takeover (`10.200.1.0/24`
    lock). `--accept-routes=true` now explicit on all `tailscale up --reset` calls.
319 - **July 2, 2026** Auto-recovery daemon documented in 6 above (`scripts/watcher.sh` + launchd plist). See `
    devref.md 15.5.1`.
320 - **July 2, 2026** Linux scripts (`scripts/toggle-linux.sh`, `recover-linux.sh`, `setup-linux.sh`) documented
    in 6.
321 - **July 1, 2026** `recover.sh --post-wake` ships; wired to watcher daemon. Triggered on every sleep/wake or
    network-state change.
322 - **June 28, 2026** 8 internal scripts moved to `scripts/`. User-facing files (`toggle.sh`, `recover.sh`, `
    setup.sh`, `docker-compose.yml`) stay at repo root.
323 - **June 28, 2026** All hardcoded install paths removed. `SCRIPT_DIR`-relative resolution and `
    __NULLEXIT_HOME__` placeholder used throughout.
324
325 ### Cryptographic Integrity Verification
326
327 nullexit uses a built-in cryptographic integrity checker designed to provide tamper-*evidence* against
    accidental edits, logic bugs, or non-privileged malware. During setup, a 256-bit `NULLEXIT_SEED` is
    injected into your `.env` file. The core entrypoint scripts (`toggle.sh`, `recover.sh`, etc.) are hashed
    using HMAC-SHA256, and their signatures are stored in `.signatures`.
328
329 *(Note: This is not a strict boundary against an attacker who already has write access to the repo, as they
    could simply read the seed and re-sign the scripts. It is designed as a strict footgun-catcher).*
330
331 If any script is modified without authorization, the gateway will loudly fail to boot. If you make intentional
    edits to the bash scripts, you must re-sign them by running:
332 ```bash
333 ./scripts/crypto.sh --sign
334 ```

```

### 3 agent.md

```

1 # nullexit Detailed Agent Analysis
2
3 > Complete per-file breakdown of every function, constant, duplication, and bug found.
4 > **Note (July 2026):** `sync-rules.py` and `logger.py` have been fully ported to Go (`scripts/rule-compiler/
    main.go` and `scripts/logger/main.go`) to dramatically improve performance and reduce container footprint
    via multi-stage Alpine builds. The analysis below reflects their previous Python states.
5
6 ## Important Agent Instruction
7 - **NEVER use `cat << EOF` or terminal redirections to write or modify files.** You must always use your native
    file-editing tools (`replace_file_content` or `multi_replace_file_content`). Using `cat EOF` leads to
    duplicated text, broken markdown headers, and structural corruption (which previously destroyed the section
    numbering in `devref.md`).
8 - **Always run `bash scripts/crypto.sh --sign` after making any modifications to the core bash scripts.** This
    is required to update the cryptographic HMAC-SHA256 signatures; otherwise, the startup integrity checks
    will fail and block execution.
9 - **Always run `git diff` and print the changes to the user *before* requesting or executing any git commit
    command.** The user must be shown the diff so they can review and understand the changes. Based on this
    diff, formulate a highly precise, detailed, and descriptive commit message (explaining exactly what was
    changed and why) rather than a generic summary, and present it to the user alongside the diff.
10 - **Always run `git status` before `git diff` to identify all modified and untracked files, ensuring no files (
    such as new scripts or configuration assets) are missed before formulating commit diffs and staging.**
11 - **Always use the `--restart` flag when asked to restart the toggle (e.g., run `./toggle.sh --restart`).**
12 - **NEVER execute any `sudo` commands (especially `sudo route`, `sudo dscacheutil`, or any network-modifying
    commands) without explicitly explaining to the user what the command does and why it is needed FIRST. Do
    not assume implicit permission. Wait for the user's approval before running it.**
13 - **When fixing an error and using logs to debug, always show the user the reasoning and the specific logs that
    support this theory.**

```

```

14 - **When the user tells you to "push", this means read git diffs and prepare commit messages that correspond to
    these changes (do not ignore any changes). If the changes can be atomized into multiple commits for easier
    understanding, do so.**
15 - **When the user says "latex this", it means you should run `python3 scripts/generate_tex.py` to regenerate
    the unified LaTeX document.**
16 - **NEVER apply changes to running gateway containers directly via `docker compose up` or `docker compose
    restart`.** Recreating or restarting containers (especially the `warp` container which hosts the network
    namespace) breaks the shared network stack for all other services (`adguardhome`, `tailscale`, `routing-fix`
    `), severing host connectivity and freezing macOS DNS. If you need to apply changes or rebuild containers,
    cleanly stop the gateway first using `./toggle.sh`, or trigger `recover.sh --post-wake` to rebuild and re-
    verify the network stack safely in the correct dependency order. Do NOT run `./toggle.sh --restart` or any
    network-rebuilding commands yourself if the execution could sever host network access; instead, explain the
    changes and ask the USER to run `./toggle.sh --restart` (or `./toggle.sh` stop and start) themselves to
    safely apply the modifications.
17
18 ## How to Verify Gateway is Working
19 Whenever modifications are made to the gateway scripts, routing, or containers, execute these verification
    checks in order:
20 1. **Container Health:** Run `docker compose ps` to ensure all containers (`warp`, `adguardhome`, `routing-fix`
    ``, `tailscale`) are `running` and `healthy`.
21 2. **DNS Hijacking Status:** Run `networksetup -getdnsservers "Wi-Fi"` (or appropriate network service) and
    ensure it is set exclusively to the gateway's Tailscale static IP (typically `100.100.21.8`).
22 3. **DNS Query Interception:** Run `dig @100.100.21.8 google.com` (using the gateway static IP from `.`
    gateway_ip`) to ensure AdGuard Home is active, intercepting, and resolving queries.
23 4. **Double-Tunnel Egress:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp` to verify the
    egress is routed through Cloudflare WARP (`warp=on`).
24 5. **Host Leak Monitoring:** Run `tail -n 30 output.log` and verify the background watchers (DNS + WARP) report
    healthy status without `LEAK` warnings. For a deeper on-demand audit, run `bash scripts/diagnose-host-leak
    .sh` (add `--watch` to poll, `--fix` to remediate).
25
26 **Automated:** `bash scripts/sweep.sh` runs a read-only 6-check health sweep covering the above plus the PF
    kill-switch and tailnet-wide peer reachability, writes a timestamped `sweep-<UTC>.txt` report to the repo
    root, and exits non-zero on FAIL (`--quiet` for just the verdict). It never mutates routing/PF/DNS/
    containers, so it is safe against a live gateway. See its entry in Part 3.
27
28 ---
29
30 ## Part 1: macOS Main Scripts
31
32 ### toggle.sh (2191 lines, unified macOS + Linux)
33
34 > **Refreshed 2026-07-15.** Post-unification (`$OSTYPE` switch at L20) + the July-2026 `common.sh` extraction,
    the tables below reflect the *current* file. Most former "toggle.sh functions" now live in `scripts/common.
    sh`; the host-leak-probe pair was removed. For the ordered execution flow see **"toggle.sh Execution Map &
    Critical Path"** above.
35
36 **Functions still defined in `toggle.sh`** several exist in *both* halves (Linux L63848 / macOS L8782191):
37
38 | Function | Linux L | macOS L | Purpose |
39 |-----|-----|-----|-----|
40 | `_log_exit_breadcrumb` | 30 | 885 | Write lifecycle-phase breadcrumb on exit (post-mortem trail) |
41 | `_get_free_port` | 67 | 944 | Pick a random unused port (self- + kernel-collision checked) |
42 | `_run_gui_cmd` | 1033 | | Run a command as the logged-in console GUI user |
43 | `_write_gateway_active_marker` | 1053 | | Write `/tmp/nullexit-gateway-active.marker` (macOS only) |
44 | `_clear_gateway_active_marker` | 1059 | | Remove active marker + `TUNNEL_FAILED_CLOSED.marker` |
45 | `_start_sleep_prevention` | 145 | 1091 | Start `caffeinate` w/ revert trap |
46 | `_stop_sleep_prevention` | 188 | | Stop `caffeinate` via PID file |
47 | `_stop_dns_watcher` | 202 | | Stop DNS watcher (macOS uses `stop_pidfile_daemon`) |
48 | `_start_warp_watcher` | 1164 | | Background WARP liveness monitor `recover.sh` after N fails |
49 | `_stop_warp_watcher` | 1241 | | Stop WARP watcher via PID file |
50 | `_cleanup_handler` | 215 | 1258 | Fail-closed teardown trap (guarded by `SUCCESS_RUN`) |
51 | `_restart_tailscaled_daemon` | 327 | *(common)* | Restart tailscaled (Linux-half local copy; macOS uses
    common.sh) |
52 | `_disconnect_tailscale_host` | 337 | *(common)* | `tailscale down` w/ retry (Linux-half local copy; macOS
    uses common.sh) |
53 | `_setup_exit_node_routing` | 1430 | | Override default route through the Tailscale utun interface |
54 | `_cleanup_network_state` | 354 | 1464 | Nuclear cleanup: proxies, DNS cache, routes, Wi-Fi, sharing services
    |
55
56 **Delegated to `scripts/common.sh`** (moved out of toggle.sh in the July-2026 refactor; line = common.sh):
57
58 | Function | common.sh L | Function | common.sh L |
59 |-----|-----|-----|-----|
60 | `_run_with_timeout` | 58 | `_enable_socks_proxy` | 769 |
61 | `_is_gateway_active` | 168 | `_disable_socks_proxy` | 803 |
62 | `_gateway_state` | 232 | `_reset_dns` | 824 |
63 | `_get_active_service` | 284 | `_stop_local_dns_proxy` | 898 |
64 | `_restart_tailscaled_daemon` | 338 | `_start_local_dns_proxy` | 909 |
65 | `_disconnect_tailscale_host` | 352 | `_force_dns_to_gateway` | 988 |
66 | `_stop_pidfile_daemon` | 393 | `_add_warp_bypass_routes` | 484 |
67 | `_start_dns_watcher` | 715 | `_remove_warp_bypass_routes` | 551 |

```

```

68
69 **Removed:** `start_host_leak_probe()` / `stop_host_leak_probe()` and the whole `HOST_LEAK_PROBE` prober
    subsystem no longer defined or called anywhere in `toggle.sh`, `recover.sh`, or `common.sh`, and `
    HOST_LEAK_PROBE` is no longer an env var (all verified). The only surviving host-leak tool is the
    standalone, on-demand `scripts/diagnose-host-leak.sh` (there is **no** `scripts/host-leak-probe.sh`).
70
71 **Constants / anchors (current):**
72
73 | Constant | Value | Line |
74 |-----|-----|-----|
75 | `LOG_FILE` | `$PWD/output.log` | L22 (Linux) / L863 (macOS) |
76 | `LOCK_FILE` | `/tmp/nullexit-toggle.lock` | L1020 |
77 | `PID_FILE` (caffeinate) | `/tmp/nullexit-caffeinate.pid` | L142 |
78 | `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L200 |
79 | `WARP_WATCHER_PID_FILE` | `/tmp/nullexit-warp-watcher.pid` | L1144 |
80 | Gateway active marker | `/tmp/nullexit-gateway-active.marker` | L1054 (write) |
81 | TUNNEL_FAILED_CLOSED marker | `

```

```

139 stop_warp_watcher` (1576) `clear_gateway_active_marker` (1577) "STOPPED in Ns" (1584).
140 ##### Long-lived background processes spawned by START
141
142 | Process | Started | PID / reap | Role |
143 |-----|-----|-----|-----|
144 | honey-port `nc -l localhost $PORT` (disowned) | L1695 | `kill -f "nc -l localhost"` | port-scan tripwire |
145 | `caffeinate` | L2130 | `/tmp/nullexit-caffeinate.pid` | prevent system sleep |
146 | DNS watcher loop | L2134 | `/tmp/nullexit-dns-watcher.pid` | re-hijack DNS every 30s (Wi-Fi roam) |
147 | WARP watcher loop | L2138 | `/tmp/nullexit-warp-watcher.pid` | monitor warp; fire recover.sh after N fails |
148 | local DNS proxy `dns-proxy.py` | L2089 | (killed by `stop_local_dns_proxy`) | UDP:53 gateway DNS fallback |
149
150 Separately, the launchd `scripts/watcher.sh` daemon (post-wake / post-roam recovery) runs independently of
    toggle and invokes `recover.sh --post-wake`; toggle does not spawn it.
151
152 ##### `cleanup_handler` the fail-closed teardown trap
153
154 Armed on `ERR/INT/TERM/HUP` (L1317-1320), guarded by `SUCCESS_RUN`. While `SUCCESS_RUN=false` it runs: `
    reset_dns`, `disconnect_tailscale_host`, `stop_local_dns_proxy`, stop watchers, `
    clear_gateway_active_marker`, `docker compose down` (+ `colima stop` on `ERR`). This is the fail-closed
    guarantee: a half-built gateway is torn down, not left leaking.
155
156 ##### Correctness findings (verified against source)
157
158 1. STARTED in Ns (L2127) is NOT the commit point `SUCCESS_RUN=true` (L2164) is. The marker write (L2146)
    and watcher starts (L2130-2138) sit between them. Any TERM/INT including a wrapping tool's timeout
    in the L2127L2164 window fires cleanup_handler` a FULL teardown (containers + colima + host tailscale),
    leaving the marker ABSENT and the host on its raw ISP link, despite the printed success. This is exactly
    the 2026-07-15 incident. (The Linux half has the same shape but a much narrower window: STARTED L814 `
    SUCCESS_RUN=true` L820, and it maintains no marker.)
159 2. Never pipe toggle.sh` through a reader (tail`, grep`). The disowned children (honey-port,
    watchers, dns-proxy, caffeinate) inherit stdout, so the pipe never EOFs and the reader hangs long after
    toggle has exited which then trips the wrapper's timeout SIGTERM finding #1. Run detached with a file
    redirect instead: nohup ./toggle.sh --restart > /tmp/t.log 2>&1 &`, then poll the file.
160 3. Host wiring is correctly on the synchronous critical path (steps 8 & 10 precede "STARTED"), so a 
    completed START cannot leave the host unrouted/leaking the only way to that state is an interruption per
    finding #1.
161 4. Namespace ordering: step 6's docker compose up` creates warp's network namespace that tailscale`/`
    routing-fix`/`adguard` share (compose depends_on: warp healthy` enforces order). Never docker compose
    restart warp` on a live gateway it detaches the shared netns (see Agent Instruction, L16).
162
163 ##### Linux path where it diverges (L63848)
164
165 The Linux half is intent-similar but not a line-for-line mirror of macOS. Verified differences:
166
167 | Aspect | macOS (L8782191) | Linux (L63848) |
168 |-----|-----|-----|
169 | VM layer | Colima (colima_start_until_ready`, L1631) | none native Docker |
170 | Host exit-node assert | $TS_BIN set --exit-node="$TS_IP`" (L2046) | $TS_BIN up --reset --exit-node="
    $TS_IP`" (L735) |
171 | gateway-active marker | written (L2146) | not maintained (L119) gateway_state` degrades to a live
    container probe |
172 | Commit window | STARTED L2127 SUCCESS_RUN` L2164 | STARTED L814 SUCCESS_RUN` L820 |
173 | Port gen tooling | jot` / netstat` | shuf`/`$RANDOM` / ss` |
174
175 The up --reset` vs set` distinction matters: --reset` re-applies all prefs and forces a route re-eval,
    whereas set` mutates only the exit-node pref the same asymmetry recover.sh --post-wake` relies on (
    devref 15.12.21).
176
177 ---
178
179 recover.sh (566 lines, 23KB)
180
181 Functions (7 total):
182
183 | # | Function | Lines | Purpose |
184 |---|-----|-----|-----|
185 | 1 | step() | L40 | Print bold step header |
186 | 2 | ok() | L41 | Print green success message |
187 | 3 | warn() | L42 | Print yellow warning message |
188 | 4 | fail() | L43 | Print red failure message |
189 | 5 | die() | L44 | Print red error and exit |
190 | 6 | run_with_timeout() | L5674 | Run command with timeout (bash-native, no GNU coreutils) |
191 | 7 | restart_tailscaled_daemon() | L118129 | Restart tailscaled daemon via brew services |
192
193 Constants:
194
195 | Constant | Value | Line |
196 |-----|-----|-----|
197 | Color codes | RED`, GREEN`, YELLOW`, BOLD`, NC` | L37-38 |
198 | SCRIPT_DIR` | $(dirname "${BASH_SOURCE[0]}")` | L50 |

```



```

199 | `PATH` export | `/opt/homebrew/bin:/usr/local/bin:$PATH` | L77 |
200 | `PID_FILE` | `/tmp/nullexit-cafeinate.pid` | L387 |
201 | `DNS_WATCHER_PID_FILE` | `/tmp/nullexit-dns-watcher.pid` | L407 |
202 | WARP endpoints | `162.159.192.1`, `162.159.193.1` | L329-336 |
203 | TUNNEL_FAILED_CLOSED marker | `${SCRIPT_DIR}/TUNNEL_FAILED_CLOSED.marker` | L51 |
204
205 **Duplicated logic from toggle.sh:**
206 *(Note: As of July 2026, **ALL** of the below duplicated logic has been successfully extracted to `scripts/
    common.sh`!)*
207 - `run_with_timeout()` simpler subshell version vs toggle's PID-tracking version
208 - `restart_tailscaled_daemon()` near-identical (uses warn()/ok() vs echo)
209 - Active network service detection inline at L217-233 (toggle has `get_active_service()` function)
210 - ENO service resolution inline at L230-233 (same as toggle L515-516)
211 - WARP bypass routes inline at L329-337 (toggle has `add_warp_bypass_routes()`)
212 - Exit node routing setup inline at L340-345 (toggle has `setup_exit_node_routing()`)
213 - DNS cache flushing L296-302 (same as toggle L611-616)
214 - Wi-Fi power-cycling L442-461 (similar to toggle L626-637)
215 - Proxy disabling L282-288 (same as toggle L604-608)
216 - Sharing services reset L586 (same as toggle L642)
217 - PID file stop pattern L387-399, L407-419, L423-435 (same as toggle L157-167, L193-203, L312-322)
218 - KILL_SWITCH check L194 (same as toggle L141, L469)
219 - .gateway_ip reading L138-142, L209-213 (same pattern as toggle L1080)
220
221 ---
222
223 ### setup.sh (410 lines, 18KB)
224
225 **Functions (4 total):**
226
227 | # | Function | Lines | Purpose |
228 |---|-----|-----|-----|
229 | 1 | `step()` | L10 | Print bold blue step header |
230 | 2 | `ok()` | L11 | Print green success message |
231 | 3 | `warn()` | L12 | Print yellow warning message |
232 | 4 | `die()` | L13 | Print red error and exit |
233
234 **Constants:**
235
236 | Constant | Value | Line |
237 |-----|-----|-----|
238 | Color codes | `RED`, `GREEN`, `YELLOW`, `BLUE`, `BOLD`, `NC` | L7-8 |
239 | `RECOMMENDED_TS_VERSION` | `1.98.5` | L71 |
240 | `ADGUARD_PASSWORD` | `nullexit` | L215 |
241 | Default .env values | `GATEWAY_RULE_PROFILE=medium`, `GATEWAY_MSS=1120`, etc. | L240-248 |
242
243 ---
244
245 ## Part 2: Linux Scripts
246
247 ### toggle.sh Linux path (unified into the root cross-platform script)
248
249 The Linux toggle is no longer a separate file. `scripts/toggle-linux.sh` was
250 merged into the repo-root `toggle.sh`, which dispatches on `${OSTYPE}`: an early
251 `if [[ "${OSTYPE}" != darwin* ]]` block runs the self-contained Linux toggle
252 (shuf / ss / nmcli / ip / systemd) and exits before the macOS body. The shared
253 DNS/proxy/daemon primitives it used were already moved to `common.sh` in the
254 `313dc32` unification, so both branches call them by the same names.
255
256 ---
257
258 ### recover.sh Linux path (unified into the root cross-platform script)
259
260 Linux recovery is no longer a separate file. `scripts/recover-linux.sh` was
261 merged into the repo-root `recover.sh`, which dispatches on `${OSTYPE}`: the Linux
262 branch runs a self-contained teardown (resolvectl / nmcli / ip) and exits before
263 the macOS dual-mode body. All formatting/timeout helpers come from `common.sh`.
264
265 ---
266
267 ### scripts/setup-linux.sh (416 lines, 18KB)
268
269 **Functions (4 total):**
270
271 | # | Function | Line | Purpose |
272 |---|-----|-----|-----|
273 | 1 | `step()` | L10 | Print formatted step header |
274 | 2 | `ok()` | L11 | Print green success message |
275 | 3 | `warn()` | L12 | Print yellow warning message |
276 | 4 | `die()` | L13 | Print red error + exit 1 |
277
278 ***95% identical to macOS setup.sh.*** Only differences:

```

```

279 1. Desktop shortcuts (L366-391 Linux `.desktop` files vs macOS `.app` via `osacompile`)
280 2. Final instructions text
281 3. .env template: macOS adds `GATEWAY_USE_EXIT_NODE`, `WARP_FAIL_THRESHOLD`, `KILL_SWITCH` Linux omits these
282
283 ---
284
285 ## Part 3: Utility Scripts
286
287 ### scripts/diagnose-host-leak.sh (722 lines, 38KB)
288
289 **Purpose:** Comprehensive host-routing diagnostic. 8 checks classify host IP leaks. Supports `--fix` and `--
    watch` modes.
290
291 **Duplicated logic:**
292 - `resolve_active_service()` identical pattern as toggle.sh, recover.sh, fix-docker-bridge-collision.sh
293 - `run_with_timeout()` reimplemented timeout
294 - Color codes RED/GREEN/YELLOW/BOLD/NC with tty detection
295 - WARP check via `cloudflare.com/cdn-cgi/trace`
296 - `export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"`
297 - .gateway_ip reading pattern
298
299 ### scripts/fix-docker-bridge-collision.sh (402 lines, 18KB)
300
301 **Purpose:** Fixes Docker bridge subnet collisions with `10.200.1.0/24`.
302
303 **Duplicated logic:**
304 - `resolve_active_iface()` same interface detection as diagnose-host-leak.sh
305 - Color codes, step()/ok()/warn()/fail()/die() formatting helpers
306
307 ### scripts/fix-ssh-delay.sh (63 lines, 2.7KB)
308
309 **Purpose:** One-shot fix for ~20s SSH delay. Standalone, minimal duplication (uses hardcoded / instead of
    color functions).
310
311 ### scripts/reinstall-host-tailscale.sh (740 lines, 31KB)
312
313 **Purpose:** Complete uninstall + reinstall of host Tailscale.
314
315 **Duplicated logic:**
316 - `with_timeout()` yet another bash timeout implementation (polling-based, different from diagnose and toggle
    versions)
317 - Color codes + logging functions
318 - System Extension detection logic (similar to diagnose-host-leak.sh)
319
320 ### scripts/routing-fix.sh (266 lines, 10.5KB)
321
322 **Purpose:** Runs INSIDE warp container. Maintains routing table 200, FORWARD rules, country blocking, IP
    blocklists, tunnel health checks.
323
324 **Duplicated logic:**
325 - (Fixed) WARP_ENDPOINT defaults are now centralized in common.sh
326 - WARP health check via cdn-cgi/trace (another instance)
327 - iptables dual-backend pattern (for-loop over `iptables iptables-legacy`)
328
329 ### scripts/sweep.sh (331 lines, 18KB)
330
331 **Purpose:** Read-only post-restart health sweep answers "after a toggle/restart/wake, is the gateway actually
    healthy end-to-end?" Never mutates routing, PF, DNS, or containers, so it is safe to run against a live
    gateway at any time (including from a remote session). `errexit` is deliberately OFF: probes are *expected*
    to fail on an unhealthy gateway and each exit status is recorded as the diagnostic signal.
332
333 **The 6 checks** (each recorded PASS/WARN/FAIL; exit 0 only when nothing FAILs, so it can gate CI/launchd/post-
    restart hooks):
334 1. **Containers** `docker compose ps`: warp/adguardhome/tailscale/tor/routing-fix up (warp + adguardhome
    additionally `healthy`).
335 2. **Double-tunnel/leak** container WARP egress IP == host egress IP, both `warp=on` (gluetun ships wget not
    curl, so the in-container trace falls back to wget).
336 3. **Hostcontainer path** the gateway peer line in `tailscale status` shows `direct <ip:port>`; WARNs on a
    DERP relay.
337 4. **AdGuard filtering** `compiled_rules.txt` block counts present + live interception (`dig @` the gateway IP
    from `.gateway_ip`).
338 5. **PF kill-switch** `pfctl` reports Enabled AND anchor `com.apple/nullexit` carries rules (needs
    passwordless sudo, else WARN).
339 6. **Tailnet reachability** (added 2026-07-17) every peer `tailscale status` reports online (self and `offline
    peers` skipped; a `--` status means online-but-idle and is kept) must answer `tailscale ping --timeout=3s
    --c 1 --until-direct=false` (TSMP probes the WireGuard data path even on devices that drop ICMP). All pong
    PASS; some unreachable WARN; none FAIL. DERP-relayed pongs count as reachable (direct-only enforcement
    stays in check 3, which covers the gateway).
340
341 **Output:** coloured stdout summary + a timestamped plain-text report `sweep-<TIMESTAMP>.txt` in the repo root
    (one file per run, so runs can be diffed); `--quiet` prints only the verdict. Signed in `.signatures` re-

```



```

run `bash scripts/crypto.sh --sign` after editing it.
342
343 ### scripts/watcher.sh (269 lines, 13KB)
344
345 **Purpose:** Long-running launchd daemon for post-wake/post-rom recovery.
346
347 **Key function `detect_lan_p2p_mode()`:**
348 Added July 10, 2026. Called on startup and on every network state change (Listener 2 NET: events). Determines
whether the current Wi-Fi network safely allows direct Tailscale LAN P2P connections:
349 1. **WPA2-Enterprise check:** `system_profiler SPAirPortDataType` reads the `Security:` field for the current
association. Enterprise = 802.1x = always AP-isolated sets `allow=false`.
350 2. **AP isolation probe:** `route -n get 1.1.1.1` finds default gateway IP, then `ping -c 3 -t 1 -q "$gw"` if
0 responses on a non-empty network, AP isolation is active sets `allow=false`.
351 3. **Output:** Writes `true` or `false` to `.lan_p2p_detected` in the repo root.
352 4. **Consumer:** `routing-fix.sh` reads this file every 30 seconds and enforces/relaxes the RFC1918 DROP rule
accordingly no restart required.
353
354 **Duplicated logic:**
355 - PATH export (similar to toggle.sh/recover.sh)
356 - Marker file pattern (`/tmp/nullexit-gateway-active.marker`)
357
358 ### scripts/dns-proxy.py (92 lines, 2.9KB) Clean, standalone
359
360 ### scripts/logger.py (144 lines, 5.9KB) Clean, standalone
361
362 ---
363
364 ## Part 4: Cross-Cutting Issues
365
366 ### Functions Identical Across macOS Linux
367
368 | Function | macOS toggle | macOS recover | macOS setup | Linux toggle | Linux recover | Linux setup |
369 |-----|-----|-----|-----|-----|-----|-----|
370 | `run_with_timeout()` | | | identical | identical | | |
371 | `stop_local_dns_proxy()` | | | identical | | | |
372 | `start_local_dns_proxy()` | | | ~95% similar | | | |
373 | `is_gateway_active()` | | | ~80% similar | | | |
374 | `step/ok/warn/die` | | | identical | identical | | |
375
376 ### Platform-Adapted Functions (same logic, different OS commands)
377
378 | Function | Linux Command | macOS Command |
379 |-----|-----|-----|
380 | `get_active_service()` | `ip route get 1.1.1.1 \| awk '{print $5}'` | `route get default` + `networksetup -
listnetworkserviceorder` |
381 | `reset_dns()` | `resolvectl revert` | `networksetup -setdnsservers` |
382 | `cleanup_network_state()` | `ip link set dev down/up` | `networksetup -setairportpower off/on` + proxy
disable |
383 | `force_dns_to_gateway()` | `resolvectl dns` + `resolvectl domain -ts.net` | `networksetup -setdnsservers` +
`-setsearchdomains ts.net` |
384 | `start_sleep_prevention()` | `systemd-inhibit --what=sleep` | `caffeinate -i` |
385 | `enable/disable_socks_proxy()` | Stub (no-op on Linux) | `networksetup -setsocksfirewallproxy` |
386 | DNS cache flush | `resolvectl flush-caches` | `dscacheutil -flushcache` + `killall -HUP mDNSResponder` |
387
388 ### Features Missing From Linux Scripts
389
390 These exist in `toggle.sh`'s macOS branch but NOT its Linux branch:
391
392 1. `write_gateway_active_marker()` / `clear_gateway_active_marker()`
393 2. `restart_tailscaled_daemon()`
394 3. `start_warp_watcher()` / `stop_warp_watcher()` WARP liveness monitor
395 5. `add_warp_bypass_routes()` / `remove_warp_bypass_routes()` / `setup_exit_node_routing()`
396 6. `LOG_FILE` structured logging (timestamps to output.log)
397 7. MSS clamping injection (step 4c)
398 8. Rule count reporting
399 9. `--post-wake` mode in recover
400 10. Container teardown on failure in cleanup_handler
401
402 ---
403
404
405
406 ## Part 5: Constant Duplication Map
407
408 | Value | Files |
409 |-----|-----|
410 | `162.159.192.1` (WARP endpoint) | docker-compose.yml, routing-fix.sh (now dynamically loaded via .env/common.
sh!) |
411 | `5335` (AdGuard DNS port) | docker-compose.yml, AdGuardHome.yaml, post-rules.txt |
412 | `5354` (mapped DNS port) | docker-compose.yml, dns-proxy.py |
413 | `41642` (Tailscale WireGuard) | docker-compose.yml, routing-fix.sh, post-rules.txt |

```

```

414 | `1080` (SOCKS5 port) | docker-compose.yml, toggle.sh |
415 | `1120` (MSS clamp) | post-rules.txt, .env |
416 | `admin:nullexit` (AdGuard creds) | AdGuardHome.yaml (bcrypt) |
417 | `10.200.1.0/24` (Docker subnet) | (Fixed) no longer duplicated |
418 | `America/NewYork` (timezone) | docker-compose.yml (4 services) |
419 | `/tmp/nullexit-caffeinate.pid` | (Fixed) centralized in common.sh |
420 | `/tmp/nullexit-dns-watcher.pid` | (Fixed) centralized in common.sh |
421 | `/opt/homebrew/bin:...` (PATH) | (Fixed) centralized in common.sh |
422
423 ---
424
425 ## Part 6: Refactoring Outcome (Implemented July 2026)
426
427 **Decision: "Scripts-only" architecture**
428 Instead of introducing a new `lib/` directory, all shared code was moved into the existing `scripts/` directory
  to keep the project hierarchy flat and simple.
429
430 ### 1. `scripts/common.sh` (288 lines, 10KB)
431 Extracts the fundamental primitives and networking logic used across orchestrator scripts:
432 - **Formatters:** `step()`, `ok()`, `warn()`, `fail()`, `die()`
433 - **Timeout Wrapper:** `run_with_timeout()` (Canonical version with PID tracking, graceful SIGTERM, and SIGKILL
  fallback)
434 - **Daemon Lifecycle:** `restart_tailscaled_daemon()`, `stop_pidfile_daemon()` (collapsed caffeinate, DNS
  watcher, and WARP watcher teardown logic)
435 - **Network Interface/Routing:** `get_active_service()`, `get_en0_service()`, `bounce_wifi_interfaces()`,
  `add_warp_bypass_routes()`, `remove_warp_bypass_routes()`
436 - **Proxy/DNS Cleanup:** `disable_all_proxies()`, `flush_dns_cache()`, `reset_sharing_services()`
437 - **Configuration Helpers:** `read_adguard_ip()`, `is_kill_switch_enabled()`, `get_warp_endpoint_1()`,
  `get_warp_endpoint_2()` (which centralizes the hardcoded 162.159.192.1 / .193.1 into .env logic)
438
439 ### 2. `scripts/setup-common.sh` (~350 lines)
440 Extracts the heavy, duplicated installation logic shared between `setup.sh` (macOS) and `scripts/setup-linux.sh`
  :
441 - Docker and Colima prerequisite checks
442 - `wgcf` binary download and WARP key generation
443 - `.env` configuration file generation
444 - Interactive Tailscale Auth Key and AdGuard Rule Profile prompts
445
446 ### Sourcing Pattern
447 - macOS root scripts (`toggle.sh`, `recover.sh`, `setup.sh`) source via:
448   `source "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/common.sh"`
449 - Scripts under `scripts/` (`watcher.sh`, `diagnose-host-leak.sh`, etc.) source via:
450   `source "$(dirname "${BASH_SOURCE[0]}")/common.sh"`
451
452 ## Agent Verbs / Protocols
453
454 ### `/sweep` (or `sweep the gateway`)
455 When the user invokes this verb, the agent MUST perform a full end-to-end connectivity, health, security, and
  performance audit of the gateway:
456 1. **WARP Validation:** Run `curl -s https://www.cloudflare.com/cdn-cgi/trace | grep warp=` (must equal `on`).
457 2. **Tor Proxy & Control Validation:**
458   - Source the dynamically generated ports from `/tmp/nullexit-ports.env` (or identify them using `docker
  compose ps`).
459   - **SOCKS5 Check:** Run `curl -s --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.org/
  api/ip` to confirm the SOCKS5 proxy successfully connects and routes through the Tor network.
460   - **Control Port Check:** Query the Tor Control Port using netcat: `echo "PROTOCOLINFO" | nc 127.0.0.1:
  $TOR_CONTROL_PORT`. Verify it returns a standard `250-PROTOCOLINFO` response, confirming the Control Port
  is active and communicating.
461   - **Circuit & Node Lookup:** Query the Tor Control Port for the active path (`GETINFO circuit-status`). For
  each active built circuit, parse the relay fingerprints, lookup their IP addresses (`GETINFO ns/id/<
  fingerprint>`), and resolve their country codes using Tor's local GeoIP database (`GETINFO ip-to-country/<
  IP>`). Include the list of active circuits, showing the role (Guard/Middle/Exit), nickname, IP, and country
  of each hop in the final sweep report.
462   - **Transparent Onion Validation:** Run `nslookup duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.
  onion` to verify it resolves to an IP within `198.18.0.0/15` benchmark range. Verify transparent dark web
  routing by running `curl -sI -H "Host: duckduckgogg42xjoc72x3sjasowoarfbgcmvfimaftt6twagswzczad.onion" http
  ://<resolved-ip>/` (should return HTTP 301/200).
463 3. **DNS Leak Validation:**
464   - Audit the upstream resolver currently utilized by the host. Run `curl -s https://edns.ip-api.com/json` to
  fetch details of the DNS resolver processing the client's requests.
465   - Verify that the resolver IP and organization returned belong to Cloudflare (e.g., AS13335) or match the
  WARP tunnel's exit range, and that **no** DNS traffic leaks to the user's raw physical ISP DNS.
466 4. **Performance Benchmark Audit:**
467   - Run the python benchmark script `python3 benchmarks/test_load.py`.
468   - Report the overall load speed and the ratio of successfully fetched vs. blocked/failed subresources for
  the targeted domains to ensure ad-blocking and MTU encapsulation are performing correctly.
469 5. **Tailscale Mesh Validation:**
470   - Run `ping -c 1 100.100.21.8` to ensure the gateway is reachable.
471   - Run `tailscale status` to identify all active/online nodes in the mesh.
472   - For every online peer returned in the status output, execute `tailscale ping --timeout=3s --c 1 --until-
  direct=false <peer-ip>` (TSMP probes the WireGuard data path and works even on phones that drop ICMP) to

```

```

473 verify connectivity and confirm healthy routing to all active mesh devices. `bash scripts/sweep.sh`
    automates this as its check 6/6 and reports each peer's path (direct endpoint vs `DERP(...)` and latency.
474 6. **Log Audit:** Run `tail -n 100 output.log | grep -iE "(error|fail|warn)"` to catch any underlying component
    failures or warnings.
475 ### `/latex` (or `generate latex`)
476 When the user invokes this verb, the agent MUST run `python3 -I scripts/generate_tex.py` to regenerate the
    documentation source code (`nullexit_unified.tex`). The agent MUST NOT attempt to compile the `.tex` file
    into a PDF (the user handles PDF compilation locally).

```

## 4 devref.md

```

1 # nullexit Development Reference & Resolved Issues
2
3 > **Last updated:** July 12, 2026
4 > **Purpose:** Provide any LLM or developer with complete project understanding, debugging history, and
    resolved issues so they can make informed changes without re-reading every file.
5 > **Diagrams:** See [diagrams.md](./diagrams.md) for system architecture, toggle.sh flowcharts, monitoring
    layer, traffic sequence, recover.sh decision tree, and the full failureself-healing map.
6
7 This reference is organized into four parts:
8 > - **Part I Architecture & Reference** (18)
9 > - **Part II Design Decisions & Threat Model** (914)
10 > - **Part III Incident Log & Resolved Issues** (15)
11 > - **Part IV Observations, Changelog & TODO** (1618)
12
13 ---
14
15 # Part I Architecture & Reference
16
17 ## 1. What Is nullexit?
18
19 nullexit is a **Tunnel-in-Tunnel VPN gateway** that chains **Tailscale** (mesh VPN) through **Cloudflare WARP**
    (exit VPN). It runs on a macOS host inside Docker containers managed by **Colima** (lightweight Linux VM).
    The goal: any device on the user's Tailscale mesh can use this gateway as an exit node, and all traffic
    exits through Cloudflare WARP achieving double encryption, ISP invisibility, and network-wide ad-blocking
    via **AdGuard Home**.
20
21 ### Traffic Flow
22 ```
23 Phone/Laptop  Tailscale mesh (WireGuard)  Mac host  Docker container
24   Tailscale decrypts  Gluetun re-encrypts  WARP tun0  Cloudflare  Internet
25   ```
26
27 ### SOCKS5 Failover Path (if exit node is unavailable)
28 ```
29 DNS:  Browser  host 127.0.0.1:53  Python proxy  TCP:5354  AdGuard  WARP DNS
30 TCP:  Apps  macOS SOCKS5 proxy (127.0.0.1:1080)  container  tun0  WARP  Internet
31   ```
32
33 ---
34
35 ## 2. Architecture
36
37 ### Containers (all share `warp`'s network namespace via `network_mode: service:warp`)
38 | Container | Image | Role |
39 |-----|-----|-----|
40 | `warp` | `qmcgaw/gluetun:v3.41.1` | Gluetun WireGuard client  Cloudflare WARP. Owns the network namespace.
    Strict firewall. |
41 | `tailscale` | `tailscale/tailscale:v1.98.4` | Advertises as exit node on the mesh. Also provides the built-in
    SOCKS5 proxy fallback via `TS_SOCKS5_SERVER=:1080`. |
42 | `routing-fix` | `alpine:3.20` | Sidecar that maintains routing tables + iptables rules every 30 seconds. |
43 | `rule-compiler` | `golang:1.22-alpine` / `alpine:3.20` | One-shot startup container that compiles 500k+ DNS
    block rules in <2s and immediately exits. |
44 | `adguardhome` | `adguard/adguardhome:v0.107.77` | DNS sinkhole for ads/trackers. Listens on port 5335.
    Upstream DNS: `127.0.0.1:53` (through WARP). |
45
46 ### Port Mappings (on host via Colima)
47 | Host Port | Container Port | Protocol | Purpose |
48 |-----|-----|-----|-----|
49 | 5354 | 5335 | TCP+UDP | AdGuard DNS |
50 | (Tailscale IP):3000 | 3000 | TCP | AdGuard web UI (Internal to Mesh) |
51 | 41641 | 41641 | UDP | Tailscale WireGuard direct |
52 | 1080 | 1080 | TCP | SOCKS5 proxy |
53
54 ### Host Components

```

```

55 - **Colima VM** Linux VM running Docker (600MB RAM + 400MB swap)
56 - **tailscaled** Standalone Tailscale daemon via `brew install tailscale` + `brew services start tailscale` (
  NOT the GUI `.app`)
57 - **macOS network settings** DNS, SOCKS proxy, routing all manipulated by `toggle.sh`
58
59 ---
60
61 ## 3. Key Files
62
63 | File | Lines | Purpose |
64 |-----|-----|-----|
65 | `toggle.sh` | ~2109 | **Main script (macOS + Linux).** Dispatches on `$OSTYPE` an early Linux block (shuf/ss
  /nmcli/ip, systemd) runs and exits before the macOS body (jot/netstat/networksetup, launchd/pf). Detects
  state, toggles gateway ON/OFF. Handles DNS hijacking, Tailscale exit node, SOCKS proxy, sleep prevention,
  timeouts, cleanup. |
66 | `setup.sh` | 406 | **One-time setup.** Installs deps (Docker, Tailscale, wgcf), generates WARP keys, writes
  `.env`, configures AdGuard via API, starts containers. |
67 | `docker-compose.yml` | 194 | Service definitions for all 5 containers. |
68 | `scripts/routing-fix.sh` | 201 | Maintains routing tables (table 200 for SOCKS5, table 52 for CGNAT) and
  FORWARD iptables rules in both nftables and legacy backends. Loads and enforces the IP blocklist via `ipset`
  . Runs in a 30-second loop. |
69 | `post-rules.txt` | 18 | Gluetun iptables rules loaded at container start. FORWARD accept for tailscale0, NAT
  MASQUERADE on tun0, DNS redirect to port 5335, IPv6 drop, TCP MSS clamping. |
70
71 | `scripts/dns-proxy.py` | 91 | Local DNS proxy. UDP:53 TCP:5354 with proper 2-byte length prefix. Fallback
  when exit node is unavailable. |
72 | `scripts/rule-compiler/` | ~800 | Unified threat compiler (ported to Go from Python). **DNS pipeline:**
  fetches remote blocklists, deduplicates against AdGuard native filters, optimizes subdomains (~50%
  reduction), compiles to AdGuard syntax. **IP pipeline:** fetches threat-intel feeds (Spamhaus, Feodo, ET,
  CINS), strips private/Tailscale ranges, outputs atomic `ipset` restore file. Built dynamically in Docker
  multi-stage build. |
73 | **`adguard/work/`** | | **Bind-mounted container data folder. Off-limits from outside Docker READ-ONLY-
  EXCEPT-VIA-`sync-rules`. See README 6.** |
74 | `recover.sh` | ~742 | Nuclear recovery script (**macOS + Linux** dispatches on `$OSTYPE`; the Linux branch
  runs a self-contained teardown via `resolvectl`/`nmcli`/`ip` and exits, macOS falls through to the dual-
  mode body). Gain: use `--post-wake` (macOS) for non-destructive refresh after sleep/wake or Wi-Fi roam
  keeps the gateway live while still re-hijacking DNS + refreshing Tailscale exit-node + force-recreating
  warp if unhealthy. Resets DNS on ALL services, disables proxies, flushes routes, stops containers, kills
  sleep prevention, power-cycles Wi-Fi. Cryptographically verified. |
75 | `scripts/diagnose-host-leak.sh` | 580+ | **One-shot host-routing diagnostic.** Runs 8 checks (Tailscale,
  SOCKS5, CDN-cgi/trace, IPv6 leak, default route, output.log hints, gateway IP, Tailscale System Extension
  conflict), classifies into ONE of three known leak scenarios (A=SOCKS5 fallback, B=IPv6 leak, C=route
  freeze) or OK, writes a timestamped report to `host-leak-diagnostic-<UTC>.txt`, and prints a ready-to-run
  fix on stdout. Pass `--fix` to apply the matched remediation and re-verify egress. Pass `--watch` (or `--
  watch 30`) to run a full baseline then continuously monitor warp/IPv6/default-route every N seconds,
  alerting on any state change. See 15.6.1. |
76 | `scripts/host-leak-probe.sh` | | **REMOVED (2026-07-15, verified deleted/untracked).** Was a continuous sub-
  second host-egress prober auto-launched via `HOST_LEAK_PROBE`. Its fail-closed role is now covered by the
  **PF kill-switch** (`enable_killswitch` in `common.sh`) + the in-container **WARP Watcher** (15.6.2); on-
  demand auditing is `scripts/diagnose-host-leak.sh`. See 15.6.1 (retained as historical rationale). |
77 | `scripts/fix-docker-bridge-collision.sh` | ~290 | **One-shot fix for 15.2.3 (Docker bridge subnet collision)
  .** Detects Docker's `172.17.0.0/16` bridge colliding with the host Wi-Fi subnet (the symptom that produces
  `default 172.17.0.1 UGScg en0` in `netstat -rn`), picks a non-colliding `docker.bip` from a small
  candidate set, idempotently edits `~/colima/default/colima.yaml` with a timestamped backup, restarts
  Colima, rebinds Wi-Fi DHCP, verifies the new default route is no longer Docker's bridge, and re-runs `
  toggle.sh` + `scripts/diagnose-host-leak.sh`. Supports `--dry` and `--skip-toggle`. |
78 | `scripts/watcher.sh` | 194 | **Long-running post-wake + post-roam daemon.** Launched by launchd LaunchAgent;
  subscribes to `com.apple.powermanagement` events via `log stream` and to network-state changes via `scutil
  n.watch`. Both fire `recover.sh --post-wake`. Single-instance lock + SCRIPT_DIR-relative resolve of `
  RECOVER`. See 15.5.1. |
79 | `scripts/common.sh` | ~40 | **Shared script primitives.** Centralizes formatted echo statements (`step`, `ok`
  , `warn`, `die`) and the safe background `run_with_timeout` execution wrapper. Sourced by all orchestrator
  scripts across macOS and Linux. |
80 | `scripts/setup-common.sh` | ~350 | **Shared installation logic.** Centralizes logic for verifying Docker,
  installing Tailscale/wgcf, generating `.env`, and prompting for auth keys. Sourced by `setup.sh` and `setup
  -linux.sh`. |
81 | `scripts/setup-linux.sh` | ~80 | **Linux sibling of `setup.sh`.** Sources `setup-common.sh`. Distro detection
  (apt/dnf/pacman), installs Linux-specific dependencies, creates `.desktop` shortcuts. |
82 | `scripts/Toggle-Gateway.bat` | ~300 | **Windows Toggle helper.** Moved from root to `scripts/`. Helps invoke
  the framework on Windows if applicable. |
83 | `scripts/reinstall-host-tailscale.sh` | ~740 | **Nuke-and-reinstall tool for host Tailscale.** Completely
  purges Tailscale, its settings, macOS Background Items (`.btm` database), and stale System Extensions.
  Wipes ALL Login Items via `sftool reset-login-items` (requires sudo). Useful for crisis recovery but
  destructive to other login items. |
84 | `scripts/fix-ssh-delay.sh` | ~100 | **Fixes SSH connection delays.** One-shot script to address SSH delays
  caused by DNS or routing issues, useful in crisis situations. |
85 | `scripts/logger/` | ~100 | **Internal logger piped from `scripts/routing-fix.sh`** inside the `warp`
  container. Ported to Go from Python. Built via Docker multi-stage build. |
86 | `black_list.txt` | 71 | Custom domains to block (ads, trackers, telemetry). Supports `$important` modifier. |
87 | `white_list.txt` | 222 | Domains to force-allow (YouTube, Apple services, etc.). Always wins over blocks. |
88 | `.env` | ~14 | WARP WireGuard keys, Tailscale auth key, rule profile. **Contains secrets & NULLEXIT_SEED.** |

```

```

89 | `.gateway_ip` | 1 | Static gateway Tailscale IP (fallback for dynamic resolution). |
90 | `scripts/unlock-files.sh` | ~20 | **One-shot stale-permission fix.** Uses atomic rename (`cp` + `mv`) to
    replace locked inodes (mode `000`/`0444` from old `chmod` decisions) with fresh writable ones no `chmod`
    called. |
91 | `scripts/crypto.sh` | ~50 | **Cryptographic integrity enforcement.** Uses `NULLEXIT_SEED` from `.env` to sign
    core bash scripts (HMAC-SHA256) and strictly verify them on start to prevent manipulation. Pass `--sign`
    or `--verify`. |
92 | `scripts/pf.conf` | 35 | **macOS native firewall ruleset.** Enforces the default-deny kill-switch, allows
    local LAN / Tailscale traffic, and applies TCP MSS Clamping (`max-mss 1160`) to prevent WireGuard MTU
    fragmentation. |
93
94 ---
95
96 ## 4. toggle.sh Flow
97
98 ### State Detection
99 `is_gateway_active()` returns true if EITHER containers are running OR host DNS is not `1.1.1.1`. This dual
    check prevents the script from starting when it should stop (e.g., containers crashed but DNS is still
    hijacked).
100
101 ### START Path (gateway OFF ON)
102 1. **Reset DNS to 1.1.1.1** Prevents deadlocks during startup
103 2. **Disconnect host Tailscale** `tailscale down` (prevents exit-node routing during container boot)
104 3. **Compile DNS rules** `docker compose run --rm rule-compiler` (Go binary inside multi-stage Docker build)
105 4. **Boot Colima VM** (if not running) via `colima start --until_ready`, which returns as soon as Docker is
    reachable (~12s) instead of blocking on Colima's ~2-min post-provision hang (15.8.7). The networking mode
    is **gated on the LAN-P2P signal** (`lan_p2p_detected`, falling back to `TAILSCALE_ALLOW_LAN_P2P` in `.env`
    ): a trusted home network requests `--network-address --network-mode bridged` (or `shared` on WPA2-
    Enterprise), while the common AP-isolated / untrusted case boots plain fast user-mode networking (`colima
    start --memory 0.6 --vm-type vz`).
106 - **Why Bridged Networking (when enabled)?** The `--network-mode bridged` and `--network-address` flags
    force the VM to receive its own IP directly from your local router's DHCP server, placing it on the same
    physical subnet as the host. This bypasses macOS network isolation and is what lets **external LAN devices
    ** (phones, laptops on the same LAN), mDNS, and local service discovery reach across the container boundary
    . It requires the `socket_vmnet` helper and is only worthwhile on a trusted LAN, which is why it is gated
    rather than always requested. **It is *not* required for direct hostcontainer P2P** that path (the gateway
    its own Mac) works even in user-mode networking via the routing-fix `.host_ips` RETURN-rule carve-out;
    see 15.3.5 and 15.8.7.
107 - **Enterprise Wi-Fi Fallback & Dynamic Roaming:** `toggle.sh` dynamically queries `system_profiler
    SPAirPortDataType` on boot. If it detects a **WPA2-Enterprise** connection (802.1X), it forces Colima to
    fall back to `--network-mode shared`. This prevents aggressive network intrusion systems on corporate/
    university Wi-Fi switches from terminating the host's Wi-Fi connection due to "MAC spoofing" (detecting
    multiple MAC addresses originating from a single authenticated port). P2P connections are impossible on
    these networks anyway due to AP Client Isolation.
108 - **Dynamic Roam Downgrading:** `toggle.sh` writes its selected network mode to `/tmp/nullexit_colima_mode
    .txt`. When roaming between networks (e.g., Home to University), `recover.sh` actively checks this state
    against the new network's security requirement (`lan_p2p_detected`). If a mismatch is detected (e.g.,
    bridged Colima on an Enterprise network), `recover.sh` automatically forces a full `toggle.sh --restart` in
    the background to safely transition the Virtual Machine's network mode and protect the host from de-
    authentication.
109 5. **Configure VM swap** 400MB swap file inside the VM to prevent OOM
110 6. **Clean corrupted AdGuard config** Remove empty `AdGuardHome.yaml`
111 7. **Start containers** `docker compose up -d`
112 8. **Wait for gateway Tailscale** Poll `tailscale status` for "offers exit node" (up to 60s, abort at 40
    consecutive NoState)
113 9. **Resolve gateway IP** From `.gateway_ip` or `docker compose exec tailscale tailscale ip -4`
114 10. **Connect host to mesh** Verify `tailscaled` is running (auto-start if needed), then `tailscale up --reset
    --ssh=true --accept-dns=false --accept-routes=true --exit-node=` ( `--accept-routes=true` must be explicit
    `--reset` reverts it to `false` otherwise, silently preventing the default route from switching to `utun
    *`)
115 11. **Pre-flight checks** [1/3] `tailscale ping` gateway, [2/3] `dig +tcp` AdGuard DNS, [3/3] WARP container
    internet
116 12. **If all pass** Set exit node + hijack DNS to gateway IP (single server, no fallback)
117 13. **If any fail** Enable SOCKS5 proxy + local DNS proxy as fallback
118 14. **Final DNS re-force** `force_dns_to_gateway` called again after exit-node transition (tailscaled clobbers
    DNS briefly)
119 15. **Enable sleep prevention** `caffeinate -i` in background to prevent idle sleep
120
121 ### STOP Path (gateway ON OFF)
122 1. Disconnect host Tailscale (`tailscale down`)
123 2. `docker compose down -t 5`
124 3. Reset DNS to 1.1.1.1
125 4. Stop local DNS proxy
126 5. **Stop sleep prevention** Kill caffeinate process
127 6. Full network state cleanup (proxies off, DNS cache flush, route flush, Wi-Fi power cycle)
128
129 ### Safety Mechanisms
130 - `run_with_timeout()` Pure-bash watchdog for every system command (15s120s)
131 - `cleanup_handler()` Trap on ERR/INT/TERM/HUP restores DNS to 1.1.1.1, kills caffeinate, and tears down
    Tailscale
132 - `force_dns_to_gateway()` Sets DNS and VERIFIES via `networksetup -getdnsservers` (3 attempts with backoff)

```

```

133 - `SUCCESS_RUN` flag Cleanup only runs if script didn't complete successfully
134 - `start_sleep_prevention()` / `stop_sleep_prevention()` Manages `caffeinate -i` via PID file at `/tmp/
    nullexit-caffeinate.pid`
135
136 ---
137
138 ## 5. Routing & Firewall Architecture (Inside Container)
139
140 ### IP Rules (`ip rule show`)
141 | Priority | Match | Table | Purpose |
142 |-----|-----|-----|-----|
143 | 99 | `to 100.64.0.0/10` | 52 | CGNAT range Tailscale routes (injected by routing-fix) |
144 | 100 | `from 172.18.0.2` | 200 | Container's own traffic (SOCKS5 proxy) tun0 |
145 | 101 | all | 51820 | Gluetun's WireGuard fwmark tun0 |
146
147 ### Table 200 (SOCKS5 proxy traffic)
148 - `162.159.192.1 via $DOCKER_GW dev eth0` WARP endpoint exception (prevents tunnel loop)
149 - `default dev tun0` Everything else through WARP
150
151 ### iptables (TWO separate stacks!)
152 - **nftables backend** (`iptables`) Used by Gluetun's `post-rules.txt`. FORWARD policy DROP.
153 - **legacy backend** (`iptables-legacy`) Used by Tailscale's `ts-forward`. FORWARD policy ACCEPT.
154 - Both stacks apply to every packet. FORWARD RELATED,ESTABLISHED must exist in BOTH.
155
156 ### post-rules.txt (loaded by Gluetun)
157 ```
158 FORWARD: ACCEPT for tailscale0 in/out
159 INPUT: ACCEPT for tailscale0
160 NAT POSTROUTING: MASQUERADE on tun0
161 MANGLE: TCP MSS clamp to 1180
162 NAT PREROUTING: Redirect DNS (53) 5335 on tailscale0 and eth0
163 ip6tables: DROP FORWARD on tailscale0
164 ```
165
166 ---
167
168 ### 5.1 Tailscale Local Network Discovery (DERP Bypass)
169 By default, Gluetun's strict NAT swallows all of Tailscale's UDP hole-punching packets, forcing Tailscale to
    fall back to incredibly slow DERP relay servers (often adding 500ms+ latency).
170
171 To fix this, we use `FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8` in `docker-compose.yml`
    (on the `warp` container). This creates `ip rule` bypasses (Priority 99, Table 199) that route all packets
    destined for private IP ranges *outside* the WARP tunnel directly to the host's `eth0`.
172 - **172.16.0.0/12** is especially critical because many modern Wi-Fi routers (and Docker itself) assign IPs in
    this block (e.g., `172.17.x.x`).
173 - This bypass allows Tailscale's UDP packets to reach devices on the same Wi-Fi directly, establishing a fast
    peer-to-peer connection while the actual web traffic inside the tunnel remains fully routed through WARP.
174
175 ---
176
177 ### 5.2 Firewalling & Per-Device Access Control
178 Because every mesh device's traffic passes through the `warp` container's network namespace, this namespace
    serves as the ultimate choke point for network-wide firewall rules.
179
180 **Where to Put Rules**
181 The right place to add firewall rules is `scripts/routing-fix.sh`. It runs a 30-second loop inside the
    namespace and already handles re-injecting rules that Gluetun resets on reconnect. **Do not use `post-rules
    .txt` for dynamic rules**, as Gluetun flushes and resets it upon VPN reconnects.
182
183 **The Dual iptables Backend Constraint (CRITICAL)**
184 The container runs two simultaneous iptables backends: `iptables` (for the nftables backend used by Gluetun)
    and `iptables-legacy` (for Tailscale). Every rule **must** be added to both stacks, or it will silently
    fail for certain traffic.
185
186 **Avoiding Duplication**
187 Because `scripts/routing-fix.sh` runs in a loop, you must always check for a rule's existence with `-C` before
    appending with `-A`. Otherwise, your rules will duplicate every 30 seconds.
188
189 **Per-Device Identification & Filtering**
190 Each mesh device has a stable `100.x.x.x` Tailscale IP, which is visible as the `src` IP in the `FORWARD` chain
    . This is how you identify devices for filtering.
191 * **Domain-level:** AdGuard Home (port 3000, credentials `admin/nullexit`) provides a REST API that supports
    per-client blocking rules based on their Tailscale IP.
192 * **IP-level:** Use `iptables` in `scripts/routing-fix.sh` (e.g., `iptables -I FORWARD -s 100.x.x.x -d <
    blocked_ip> -j DROP`).
193 * **Country-level:** Add `ipset` to the `routing-fix` apk install step in `docker-compose.yml`, and load CIDR
    ranges from `ipdeny.com` into an ipset, then block that set in the `FORWARD` chain. See 5.3 (Geo-IP
    Blocking) for the full configuration flow.
194
195 ### 5.3 Geo-IP Blocking

```



```

196 To block traffic to and from specific countries, the system dynamically downloads country-specific IP ranges
    from [ipdeny.com](http://www.ipdeny.com/ipblocks/data/countries/) and drops them via iptables `FORWARD`
    rules.
197
198 To add a new country to the blocklist:
199 1. Find the 2-letter ISO country code (e.g., `cn` for China, `ru` for Russia, `il` for Israel, `kp` for North
    Korea).
200 2. Open your `.env` file.
201 3. Update the `BLOCKED_COUNTRIES` variable: e.g., `BLOCKED_COUNTRIES="kp il cn ru"`.
202 4. Run `toggle.sh` to restart the gateway and apply the new environment variables to the routing-fix container.
203
204 > [!NOTE]
205 > Limitations of Geo-IP Blocking: IP-based blocking is highly effective at blocking servers, direct apps,
    and raw infrastructure physically located and registered in a specific country (e.g., `www.gov.il`).
    However, it will not block high-profile websites (e.g., `jpost.com`) that route their traffic through
    global Content Delivery Networks (CDNs) like Google Cloud, Cloudflare, or Fastly. Because the CDN's IP
    addresses are registered in the US or globally, the network traffic physically never routes to the blocked
    country, allowing it to bypass the Geo-IP firewall.
206
207 ## 6. macOS Quirks to Remember
208
209 1. Colima SSH tunnel is TCP-only UDP ports (like 5354/udp) are NOT accessible from host. Use `dig +tcp` or
    the Python DNS proxy (UDPTCP converter).
210 2. Docker compose exec -T injects `\r` Always `tr -d '\r'` when capturing output for comparisons or `
    networksetup` commands.
211 3. Brew services can get stuck Fix with `sudo brew services restart tailscale`. Common state: `brew
    services list` shows `started` but `tailscale status` shows "Tailscale is stopped."
212 4. No `timeout` command macOS lacks GNU `timeout`. Use the `run_with_timeout()` bash function.
213 5. Sudo -v prompts even with NOPASSWD Use `sudo -n` (non-interactive) everywhere.
214 6. DNS is scoped to en0 macOS scutil DNS resolver is scoped to en0 (Wi-Fi). DNS changes on other
    interfaces (like USB Ethernet) are ignored unless en0's service is also updated. The script always sets DNS
    on BOTH `$ACTIVE_SERVICE` and `$ENO_SERVICE`.
215 7. Tailscaled clobbers DNS during exit-node transition `force_dns_to_gateway` is called a second time at
    the end of the ENABLE branch to counteract this.
216 8. Tailscale up --reset resets all unspecified flags to defaults Every `--reset` call MUST also pass `--
    accept-dns=false` explicitly or tailscaled re-enables DNS management.
217 9. Docker compose ps CWD pitfall `docker compose` commands require a `docker-compose.yml` in the current
    working directory. Use `docker ps` for raw container listing from any CWD; `docker compose ps` requires
    the project directory.
218
219 ---
220
221 ## 7. Environment Variables
222
223 ### `.env` File
224 | Variable | Purpose |
225 |-----|-----|
226 | `WIREGUARD_PRIVATE_KEY` | WARP WireGuard private key (from `wgcf generate`) |
227 | `WIREGUARD_PUBLIC_KEY` | WARP WireGuard public key |
228 | `WIREGUARD_ADDRESSES` | WARP WireGuard address (e.g., `172.16.0.2/32`) |
229 | `TS_AUTHKEY` | Tailscale auth key for the container |
230 | `NULLEXIT_SEED` | 256-bit seed for HMAC-SHA256 script integrity signing (generated by `setup.sh`) |
231 | `GATEWAY_RULE_PROFILE` | Rule compilation tier: `light`, `medium`, `heavy` |
232 | `GATEWAY_BYPASS_PING` | (Optional) `true` to proceed even if pre-flight checks fail |
233 | `GATEWAY_USE_EXIT_NODE` | (Optional) `false` to skip exit node, DNS-only mode |
234 | `GATEWAY_HIJACK_HOST` | (Optional) `false` to skip DNS hijacking on the host (VPN/adblocking for Tailscale
    peers only) |
235 | `GATEWAY_MSS` | (Optional) TCP MSS clamp value (default 1120); 1180 for speed on healthy paths |
236 | `WARP_FAIL_THRESHOLD` | (Optional) Consecutive `warp=off` polls before auto-shutdown (default 6 = 30s; 3 = 15
    s) |
237 | `WARP_ENDPOINT_1` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 1 (default `162.159.192.1`) |
238 | `WARP_ENDPOINT_2` | (Optional) Override Cloudflare WARP WireGuard endpoint IP 2 (default `162.159.193.1`) |
239 | `KILL_SWITCH` | (Optional) `true` to enforce strict PF lock drops all host traffic if VPN fails. Breaks SSH
    if VPN dies. |
240 | `STOP_COLIMA_ON_EXIT` | (Optional) `true` to fully shut down the Colima VM on toggle-off (saves battery on
    dedicated hosts) |
241 | `ADGUARD_USER` | (Optional) AdGuard Home web UI username (default: `admin`) |
242 | `ADGUARD_PASSWORD` | (Optional) Shared password for AdGuard Home and Tor ControlPort (default: `nullexit`) |
243 | `BLOCKED_COUNTRIES` | Space-separated 2-letter ISO codes to block via ipdeny.com CIDR ranges (e.g. `kp il cn
    ru`) |
244
245 ---
246
247 ## 8. Diagnostic Commands
248
249 ### Container Health
250 ```bash
251 docker ps --format '{{.Names}} {{.Status}}'
252 docker compose exec -T tailscale tailscale status
253 docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace
254 ```

```

```

255
256 ### SOCKS5 Proxy
257 ```bash
258 curl --socks5-hostname 127.0.0.1:1080 --max-time 5 -s https://www.cloudflare.com/cdn-cgi/trace
259 ```
260
261 ### AdGuard DNS
262 ```bash
263 dig +tcp @127.0.0.1 -p 5354 google.com +short +timeout=5
264 ```
265
266 ### Firewall & Routing
267 ```bash
268 # nftables FORWARD chain
269 docker compose exec -T warp iptables -L FORWARD -v --line-numbers
270 # Legacy FORWARD chain
271 docker compose exec -T warp iptables-legacy -L FORWARD -v --line-numbers
272 # IP rules + routing tables
273 docker compose exec -T warp ip rule show
274 docker compose exec -T warp ip route show table 199
275 docker compose exec -T warp ip route show table 200
276 ```
277
278 ### Host macOS
279 ```bash
280 tailscale status
281 brew services list | grep tailscale
282 networksetup -getdnsservers Wi-Fi
283 networksetup -getsocksfirewallproxy Wi-Fi
284 sysctl net.inet.ip.forwarding
285 ```
286
287 ---
288
289 # Part II Design Decisions & Threat Model
290
291 ## 9. Privacy Architecture & Threat Model
292
293 ### Layered Defense
294 The gateway stacks four layers targeting different attack surfaces:
295
296 - **Layer 1 ISP-Level Surveillance (WARP):** Your ISP sees only an encrypted WireGuard tunnel to a Cloudflare
297   IP. In the US, ISPs can legally sell your browsing metadata and are subject to secret NSLs (National
298   Security Letters). WARP removes their visibility entirely.
299 - **Layer 2 DNS Tracking (AdGuard Home):** DNS is the phone book of the internet. By default, queries go to
300   your ISP's resolver, giving them a complete log of your traffic. AdGuard intercepts all DNS from mesh
301   devices, blocks trackers, and resolves queries through the WARP tunnel.
302 - **Layer 3 Device Identity & IP Exposure (Tailscale):** On untrusted networks (coffee shops, public Wi-Fi),
303   your device's IP is exposed. Tailscale creates an encrypted WireGuard mesh between your devices, routing
304   all traffic to your gateway exit node.
305 - **Layer 4 Kernel-Level IP Blocking (ipset/iptables):** Sophisticated malware bypasses DNS entirely with
306   hardcoded IPs. The `rule-compiler` fetches ~16,700 threat-intelligence IPs/CIDRs (Spamhaus, Feodo Tracker,
307   Emerging Threats, CINS) and enforces them as `DROP` rules in the kernel `FORWARD` chain blocking both
308   outbound C2 connections and inbound attack traffic.
309
310 ### The Documented Threat: Why This Matters
311 This architecture is calibrated against documented mass-surveillance programs revealed in the Snowden
312 disclosures:
313 - **PRISM:** Bulk collection of communication content from major tech providers.
314 - **UPSTREAM:** Tapping physical backbone fiber, collecting metadata and content in bulk.
315 - **XKeyscore:** Retroactive search of unencrypted traffic.
316 - **MUSCULAR:** Infiltration of unencrypted internal datacenter interconnect links.
317
318 Passive bulk collection is not paranoia it is a documented reality. By double-encrypting and routing traffic,
319 nullexit prevents your ISP from harvesting your metadata.
320
321 ### Trust Assumptions & Shifting
322 Every component shifts trust rather than eliminating it:
323 1. **Physical Device Integrity:** You trust your physical devices are not compromised.
324 2. **Cloudflare Integrity:** You trust Cloudflare not to correlate your WARP tunnel with exit traffic.
325 3. **Tailscale Integrity:** You trust Tailscale's coordination server not to MITM WireGuard keys.
326
327 The goal: no single provider has a complete picture. Your ISP sees an encrypted tunnel. WARP sees traffic with
328 no account identity. Tailscale sees node topology but not content.
329
330 ## 9.1 Exit Node IP Rotation (Design Decision)
331
332 A common question for privacy architectures is whether the system should aggressively rotate its public exit IP
333 (e.g., every 5 minutes, like a Tor circuit or a proxy rotator).

```



```

322 `nullexit` uses Cloudflare WARP via Gluetun. WARP provides **anonymity in a crowd** rather than aggressive
    rotation. Thousands of users share the same Cloudflare Edge IP simultaneously, making your traffic
    indistinguishable from the herd. The IP may occasionally rotate on its own (historically surfaced as a
    ROTATE event in `output.log`), but it remains generally stable.
323
324 **Why aggressive IP rotation was intentionally excluded:**
325 1. **Broken TCP Connections:** Every rotation instantly drops all long-lived TCP sessions (SSH, large downloads
    , WebSockets, gaming, Zoom).
326 2. **CAPTCHA & 2FA Hell:** Modern web security (e.g., Cloudflare, Akamai) aggressively flags IP-hopping as bot
    behavior. Aggressive rotation guarantees the user will be bombarded with "Verify you are human" CAPTCHAs
    and "New Login Detected" 2FA emails constantly.
327 3. **WireGuard Architecture:** WireGuard is stateless and does not natively support IP hopping on the fly. To
    rotate an IP, the VPN client must actively drop the tunnel, request a new endpoint, and re-handshake. This
    introduces latency gaps where the kill-switch would repeatedly trigger and blackhole the user's internet.
328
329 **Verdict:** For a daily-driver secure gateway, shared-IP crowding (WARP) is superior to aggressive rotation.
    As an added benefit, Cloudflare WARP IPs are actually highly static (tied to the persistent WireGuard
    public key generated for your device identity). Because your IP stays static and is part of Cloudflare's
    own trusted network, it builds a high "trust score," virtually eliminating CAPTCHA loops while still hiding
    you within the massive crowd of other users in that data center. All mesh devices routing through your
    gateway will reliably share this exact same trusted IP.
330
331 *(Note: The only exception where aggressive IP rotation is genuinely required is for specialized offensive
    security tasklike high-volume web scraping or dark web research where avoiding IP bans or active tracking
    overrides the need for TCP stability.)*
332
333 ### 9.2 Cloudflare WARP Privacy & Logging
334
335 Because `nullexit` uses Cloudflare WARP (via Gluetun) as its upstream exit node, the system inherits Cloudflare
    's consumer privacy policy.
336
337 **What Cloudflare DOES NOT log (Strict No-Logs Policy):**
338 * **Browsing History:** They do not log the domains or URLs you visit.
339 * **Traffic Destinations:** They do not log the IP addresses of the servers you connect to.
340 * **Content:** They do not inspect or log the payload/data of your traffic.
341 * **Data Brokering:** They explicitly guarantee they will never sell or rent your data to advertisers.
342
343 **What Cloudflare DOES log (Minimal Telemetry):**
344 * **Data Transfer Volume:** Total bandwidth used (required to enforce 10GB/mo quotas if you are not using a
    paid WARP+ key).
345 * **Performance Metrics:** Anonymized connection speeds and latency to optimize their network routing.
346 * **Aggregate Volume:** Total traffic volumes to popular destinations in aggregate (e.g., "datacenter X sent 50
    TB to Netflix"), stripped of all user IDs.
347 * **Device ID (Public Key):** They store the persistent WireGuard public key generated by your container to
    authorize the connection.
348
349 **Privacy Verdict:**
350 Cloudflare WARP provides excellent **historical privacy**. Because they do not log your destinations, they
    cannot comply with historical subpoenas asking what websites you visited yesterday. However, it is **not
    absolute anonymity**. They still know your real ISP IP address while you are actively connected. For daily-
    driver privacy against ISPs, hackers, and corporate trackers, it is top-tier. For nation-state evasion, use
    Tor (see 11).
351
352 ### 9.3 OPSEC: Python Runtime Airgap (Isolated Mode)
353 `nullexit` explicitly relies on the host's native system Python 3 runtime for critical components like `scripts
    /dns-proxy.py`. Because this script runs with elevated privileges (`sudo -n`) to bind to the privileged UDP
    /53 socket, modifying or polluting this Python environment is strictly forbidden.
354
355 **The Zero-Dependency Rule & Isolated Mode (-I)**
356 You must **NEVER** install third-party `pip` packages (e.g., `requests`, `pyyaml`) into the system Python
    environment that `nullexit` uses.
357 - All Python scripts in this repository must be written using *only* the Python Standard Library (`socket`,
    `urllib`, `json`, `ssl`).
358 - Third-party packages introduce a massive supply-chain attack vector. A compromised `pip` package installed
    globally could allow local malware to hook into the Python runtime and exploit the `sudo` privilege of the
    DNS proxy to gain full root access.
359 - To enforce strict immunity against user-installed malicious packages, `toggle.sh` explicitly invokes the
    proxy using **Python's Isolated Mode** (`python3 -I`). This flag forces the interpreter to completely
    ignore `PYTHONPATH`, `PYTHONHOME`, and the user's `site-packages` directory, executing exclusively from the
    read-only, Apple-signed system framework.
360
361 ### 9.4 Recommended Upgrades
362
363 | Layer | Default | Upgraded |
364 |---|---|---|
365 | VPN Exit | Cloudflare WARP (US) | Mullvad (Swedish, proven no-logs, anonymous payment) |
366 | DNS Resolver | AdGuard via WARP | AdGuard via Mullvad DoH (Cloudflare-blind) |
367 | Coordination | Tailscale (US, OAuth) | Headscale (self-hosted, no third-party) |
368 | Auth | OAuth / persistent keys | Ephemeral auth keys (no identity persistence) |
369 | Content | WireGuard asymmetric | WireGuard + PSKs (coordination server blind) |
370

```

```

371 ##### Mullvad (Replacing WARP)
372 Swedish-based. Police raided their offices in 2023 and left empty-handed no logs exist to seize. Accounts are
    random numbers. Payment by cash or Monero. Replace `VPN_SERVICE_PROVIDER=custom` with `VPN_SERVICE_PROVIDER
    =mullvad` in `docker-compose.yml`.
373
374 ##### Mullvad DoH Upstreams
375 Point AdGuard's upstream resolvers to `https://adblock.dns.mullvad.net/dns-query`. Cloudflare WARP only sees
    encrypted HTTPS to Mullvad's IPs completely blind to your DNS queries.
376
377 ##### Headscale
378 Self-hosted Tailscale coordination server. Eliminates Tailscale the company, keeping all node topology under
    your control.
379
380 ##### Ephemeral Auth Keys
381 Generate in the Tailscale admin panel. Used once to authenticate; node disappears from the admin panel after
    disconnect. Breaks persistent identity linkage.
382
383 ##### Pre-Shared Keys (PSKs)
384 Add a WireGuard PSK between peer nodes. Adds a symmetric encryption layer that Tailscale's coordination server
    has no knowledge of, blinding it from reading transit content.
385
386 ### 9.5 Structural Limitations
387 - **Traffic Analysis / Metadata:** Passive fiber tapping (UPSTREAM) can observe timestamps, data volumes, and
    IP ranges without reading content. Defeating traffic analysis requires mixing networks (Tor) or continuous
    dummy packet padding both heavily degrade performance.
388 - **Targeted State-Level Adversaries:** Consumer-grade VPNs are not a complete shield against a nation-state
    actively targeting you. This stack is designed to defeat bulk passive surveillance and corporate dragnet
    collection.
389
390 ---
391
392 ## 10. Per-Device Access Control
393
394 Because every mesh device routes internet-bound traffic through the `warp` container's `FORWARD` chain, they
    all appear with their stable `100.x.x.x` Tailscale IP as the `src` address. This gives two powerful
    enforcement surfaces:
395
396 ### IP & Port Level (`scripts/routing-fix.sh`)
397 Write raw `iptables` rules targeting specific devices. The 30-second idempotent loop auto-survives Gluetun
    reconnects:
398
399 ```bash
400 # Block a specific device from a target subnet
401 iptables -I FORWARD -s 100.x.x.x -d 203.0.113.0/24 -j DROP
402
403 # Restrict a device to only HTTP/HTTPS
404 iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 3000 -j DROP
405 iptables -I FORWARD -s 100.x.x.x -p tcp ! --dport 443 -j DROP
406
407 # Block a device from internet egress entirely (mesh only)
408 iptables -I FORWARD -s 100.x.x.x -j DROP
409
410 # Time-based (e.g., cut off 10 PM 7 AM)
411 iptables -I FORWARD -s 100.x.x.x -m time --timestart 22:00 --timestop 07:00 -j DROP
412 ```
413
414 ### DNS Level (AdGuard REST API)
415 AdGuard Home supports per-client rules keyed by Tailscale IP:
416
417 ```bash
418 curl -X POST http://localhost:80/control/clients/add \
419   -H "Content-Type: application/json" \
420   -d '{
421     "name": "kids-ipad",
422     "ids": ["100.x.x.x"],
423     "blocked_services": ["youtube", "tiktok", "instagram"],
424     "safesearch_enabled": true
425   }'
426 ```
427
428 ---
429
430 ## 11. Tor & OPSEC Architecture
431
432 ### 11.1 Tor-over-WARP Topology
433
434 When extreme anonymity is required, integrating the Tor network into `nullexit` provides a powerful Tor-over-
    VPN topology:
435 * **The University/ISP** sees only an encrypted Cloudflare WireGuard tunnel, rendering them completely blind to
    the fact that you are using Tor (effectively bypassing Enterprise firewall Tor blocks).

```

```

436 * **Cloudflare** sees an encrypted TCP stream heading to a Tor Guard Node. They cannot see your DNS lookups,
    the destination website, or the content of your traffic.
437 * **The Tor Exit Node** sees your traffic, but not your real IP (only the Cloudflare IP).
438
439 ##### The "Transparent Proxy" Trap (What NOT to do)
440 It is tempting to build a system-wide "Transparent Tor Proxy" that forces all host traffic (e.g., standard
    Chrome/Safari browsers, background iCloud syncs) through Tor automatically. **This is a massive OPSEC trap
    and destroys anonymity.**
441 Modern operating systems and browsers will instantly leak your identity through background syncing (e.g., Apple
    Mail, Google Drive) and browser fingerprinting (canvas fingerprints, WebRTC, screen resolution). The Tor
    network protects your *IP*, but if your payload contains identifying cookies or unique fingerprints, the
    Tor Exit Node (which could be run by malicious actors) will instantly tie your session to your real
    identity.
442
443 ##### The nullexit Implementation: Isolated Procedural Proxy
444 To safely integrate Tor without triggering the transparent proxy trap, `nullexit` implements an **Isolated Tor
    SOCKS5 Proxy** with **Procedurally Generated Ports**:
445 1. **Isolated Container:** A lightweight `tor` Docker container runs securely inside the `warp` network
    namespace. It does not hijack system traffic.
446 2. **Opt-In Usage:** You must consciously configure a highly-hardened browser (like the official Tor Browser or
    a dedicated Firefox profile) to use the proxy. This prevents accidental background sync leaks.
447 3. **Procedurally Generated Ports:** Rather than hardcoding the standard proxy ports (like `127.0.0.1:9050` or
    `1080`), `toggle.sh` randomizes all internal proxy ports on every boot into the ephemeral range
    (10000-65000) and silently writes them to `/tmp/nullexit-ports.env`. This completely defeats static
    fingerprinting by malware.
448 4. **Kernel-Level Collision Avoidance:** To prevent the randomizer from picking a port already in use by a root
    daemon or Apple service (which would crash the gateway), the script directly queries the macOS kernel
    socket table (`netstat -an -p tcp`) to verify the port is absolutely free before assigning it.
449 5. **The Honey-Port Tripwire:** While procedural ports defeat static fingerprinting, advanced malware might
    attempt a full sequential port scan of `localhost` to find the hidden proxy. To counter this, `toggle.sh`
    spawns a "Honey-Port" (a fake random port with a netcat listener attached). If any local application
    attempts to connect to the Honey-Port, it instantly logs a critical security alert to `output.log` and
    fires a macOS desktop notification, serving as an early-warning Intrusion Detection System (IDS) for local
    malware.
450
451 > **Threat Model note:** The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). Port
    randomization and the Honey-Port only defeat blind, generic port-scanners they do NOT protect against a
    targeted local process that reads `/tmp/nullexit-ports.env`, greps process memory, or runs `lsof -i`. The
    proper mitigation is a host-level loopback-filtering firewall (Lulu/Little Snitch). See 15.9 (Threat Model:
    Local Proxy Discovery) for the full analysis and the rogue-binary verification test.
452
453 ### 11.2 Hardening: Tor Bridges and obfs4 Pluggable Transports
454 For users operating under severe Deep Packet Inspection (DPI), `nullexit` supports configuring Tor to connect
    exclusively through private bridges using the `obfs4` pluggable transport.
455
456 When `TOR_USE_BRIDGES=true` is set:
457 - **What it does:** It hides the entry IP from being matched against the public Tor consensus list, and
    disguises the traffic's protocol fingerprint from the WARP exit provider (Cloudflare).
458 - **What it does NOT do:** It does *not* add any extra layers of anonymity beyond what Tor already provides. It
    is purely a censorship-circumvention and obfuscation mechanism.
459 - **Limitations:** `obfs4` is highly effective against passive DPI, but it is not guaranteed to defeat active-
    probing-resistant detection by a highly sophisticated state-level adversary.
460
461 > **Architecture Rule:** The `nullexit` gateway deliberately does NOT bundle, hardcode, or automatically fetch
    bridge lines. Bridge lines are highly sensitive and expire. Users must obtain their own bridge lines from
    `https://bridges.torproject.org` and manually populate the `tor-bridges.txt` file. If bridges are enabled
    but the file is missing or empty, the Tor container will intentionally crash and fail to start rather than
    silently falling back to public guard relays.
462
463 **The Bootstrapping Paradox (Out-of-Band Key Exchange)**
464 Automating the bridge acquisition process (e.g., scripting the gateway to log into the Telegram `@GetBridgesBot`
    via MTProto API) introduces catastrophic OPSEC flaws:
465 1. **Identity Leaks:** Baking a Telegram session into the container explicitly links the "anonymous" gateway to
    a real-world physical phone number.
466 2. **The Chicken & Egg Deadlock:** In heavily censored regimes, Telegram itself is usually blocked. If the
    container requires Telegram to fetch bridges to bypass the firewall, it will fail because it cannot bypass
    the firewall to reach Telegram.
467 3. **Bridge Exhaustion & CAPTCHAs:** Automated scraping behaves identically to state-sponsored scrapers,
    triggering Tor's anti-bot CAPTCHAs which leads to silent proxy failures and exhausts the limited public
    bridge pool.
468
469 Instead, users should rely on a less secure layer to manually reach Telegram, obtain the bridges, and populate
    `tor-bridges.txt`. This can be done via the standard `nullexit` WARP tunnel, a mobile phone on cellular
    data, Mullvad VPN, or a secure remote VPS. *(Note: We plan on fully integrating Mullvad and VPS
    architectures directly into `nullexit` in the future, but have not had the time to build it yet).* This
    intentionally "air-gaps" the cryptographic key exchange from the proxy infrastructure, leaving zero API
    tokens and zero network traces for an adversary to exploit.
470
471 ### 11.3 Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification
472
473 ##### Context

```

```

474 Unlike the host's SOCKS5 fallback path (whose kill-switch interaction is documented in 15.7), the `tor`
    container's connection status and fail-closed behavior under VPN tunnel drops is not governed by macOS `pf`
    rules. Instead, it relies on network namespace inheritance inside Docker Compose.
475
476 ##### Mechanics
477 1. **Shared Network Namespace:** The `tor` service is configured with `network_mode: service:warp` in `docker-
    compose.yml`. This shares the network namespace of the `warp` (Gluetun) container.
478 2. **Inherited Egress Rules:** All socket communication within the namespace is bound to the same networking
    stack. Gluetun manages this stack by redirecting traffic through `tun0` (the WireGuard/WARP tunnel) and
    applying `iptables` rules that explicitly drop outbound traffic originating on `eth0` (the physical/bridge
    interface), except for permitted local subnets or the VPN server's endpoint.
479 3. **Tunnel Fail-Closed:** If the WARP connection drops or the `tun0` interface is brought down, the routing
    table entries for `tun0` become invalid or disappear. Because the `iptables` rules in the shared namespace
    continue to drop any outgoing internet traffic attempting to exit via `eth0`, the `tor` container cannot
    leak raw traffic to the host's underlying networks. Egress fails closed.
480
481 ##### Verification / Manual Test Case
482 To manually verify that the Tor container fails closed and does not leak traffic when the tunnel dies:
483
484 1. **Verify active path:** Ensure the gateway is active (`toggle.sh start`) and fetch the Tor port (`
    $TOR SOCKS_PORT` in `/tmp/nullexit-ports.env`). Verify Tor traffic routes correctly:
    ```bash
    curl --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.org/api/ip
    ```
    *(This should output a Tor exit node IP).*
485
486 2. **Simulate tunnel drop:** Bring down the virtual interface inside the shared network namespace:
    ```bash
    docker compose exec warp ip link set dev tun0 down
    ```
487
488 3. **Re-run the egress check:**
    ```bash
    curl --socks5-hostname 127.0.0.1:$TOR SOCKS_PORT https://check.torproject.org/api/ip
    ```
489
490 * **Expected Result:** The request must hang and eventually time out, or immediately fail with a proxy error
    (e.g., `curl: (97) SOCKS5: connection failed` or `curl: (7) Failed to connect`).
491
492 * **Leaked Egress Check:** Check the host's Wi-Fi router or firewall logs. No outbound packets from the Tor
    container should route directly over the bridge network to the internet.
493
494
495 ##### 11.4 Transparent .onion Routing via RFC 2544 Subnet (198.18.0.0/15)
496
497 ##### Context
498 To enable seamless dark web browsing across the entire mesh, the gateway requires a mechanism to dynamically
    map `.onion` domains to IP addresses that the host operating system and network layers can route.
499
500 ##### IP Address Conflict & Solution
501 1. **Initial Conflict:** Initially, the Tor virtual address network was configured to map `.onion` sites within
    the `10.192.0.0/10` range. However, the Colima Docker daemon dynamically allocated `10.200.1.0/24` to the
    containers. Because this Docker subnet mathematically falls within the `10.192.0.0/10` block, a routing
    loop occurred: all host packets destined for the Docker containers on ports like `9050` or `9051` were
    intercepted by the transparent proxy NAT rules and dropped, causing proxy failures.
502
503 2. **Tailscale Private IP Exclusion:** Shifting the subnet to another private range like `10.123.0.0/16`
    resolved the port conflict, but led to connection timeouts. On macOS, Tailscale's Exit Node routing
    actively excludes RFC 1918 private subnets (e.g. `10.0.0.0/8`, `172.16.0.0/12`, `192.168.0.0/16`) to
    maintain accessibility to local LAN resources (like printers/routers). Thus, packets targeting `10.123.x.x`
    were dropped locally by Tailscale before reaching the gateway tunnel.
504
505 3. **The Benchmarking Subnet Solution:** To bypass these private IP exclusions without manual static routing
    tables, Tor's virtual network was changed to `198.18.0.0/15` (reserved by RFC 2544 for benchmarking).
    Since this is not a private RFC 1918 range, Tailscale exit-node routing automatically encapsulates and
    routes all traffic destined for `198.18.0.0/15` through the tunnel to the gateway.
506
507 4. **Transparent NAT Redirection:** The `routing-fix` sidecar applies NAT rules in the shared network namespace
    :
    ```bash
    iptables -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040
    ```
508
509 This intercepts all incoming TCP traffic targeting the mapped `.onion` range and transparently proxies it
    through Tor's transparent port (`9040`).
510
511 ##### Exit Node Exclusions
512 The gateway supports the ability to exclude specific countries from being used as exit nodes (e.g. to avoid
    certain jurisdictions or nodes with poor network latency). This is configured globally via the `
    TOR_EXCLUDE_EXIT_NODES` environment variable in `.env`, which gets dynamically appended to Tor's `
    ExcludeExitNodes` and strictly enforced with `StrictNodes 1`.
513
514 ---
515
516 ## 12. Censorship-Resistant Transport & Egress Compartmentalization
517
518 ### 12.1 Why this is needed

```

525 WARP can be blocked from deployment countries with strict internet controls. The most likely cause isn't that " encrypted traffic is blocked" in general it's that WireGuard has a recognizable handshake signature, and ` WIREGUARD\_ENDPOINT\_IP=162.159.192.1` on `WIREGUARD\_ENDPOINT\_PORT=2408` is a fixed, easily-enumerable target (Cloudflare's known WARP range). That combination is exactly what IP/ASN-based and DPI-based blocking is good at catching.

526

527 **### 12.2 Important correction: Gluetun's `SHADOWSOCKS=on` is the wrong direction**

528 Gluetun's built-in Shadowsocks option runs a Shadowsocks **\*\*server\*\*** inside the already-connected VPN tunnel, so other LAN devices can reach out through Gluetun's **\*existing\*** connection via the SS protocol. It is not a way to make Gluetun itself connect **\*\*outbound\*\*** through a Shadowsocks server Gluetun only speaks OpenVPN or WireGuard as its actual transport. There is no `VPN\_TYPE=shadowsocks`. Confirmed against Gluetun's own maintainer/community discussion: people have asked for a "Shadowsocks-client" mode specifically to chain through Gluetun, and the standing answer is "run a separate Shadowsocks client container."

529

530 **### 12.3 Diagnose before building anything**

531 The right fix depends entirely on what's actually being blocked:

532 - Point Gluetun at a **\*\*different provider/IP/ASN\*\*** (any other Gluetun-supported WireGuard/OpenVPN provider, or a self-hosted WireGuard server) and test. If that gets through, the block is Cloudflare/WARP-specific, not a general WireGuard block **\*\*no new component needed\*\***, just swap `VPN\_SERVICE\_PROVIDER` / the endpoint.

533 - If WireGuard fails regardless of endpoint, the block is protocol-level (DPI fingerprinting the handshake) proceed to obfuscation.

534

535 **### 12.4 Path A: Swap WARP for a different Gluetun-native transport**

536 Simplest fix, only applies if 12.3 shows the block is Cloudflare-specific. Change `VPN\_SERVICE\_PROVIDER` / `VPN\_TYPE` / endpoint in `docker-compose.yml` and in the now-centralized `WARP\_ENDPOINT\_\*` values in `.env`. Nothing else in the stack needs to change.

537

538 **### 12.5 Path B: Add an obfuscation hop**

539 Needed if WireGuard itself is being fingerprinted/blocked. Two ways to slot it in:

540

541 **\*\*B1 Wrap (recommended):\*\*** Run a Shadowsocks (or obfs4/v2ray) client as a new sidecar container. Point Gluetun's `WIREGUARD\_ENDPOINT\_IP` at `127.0.0.1:<local-port>` instead of Cloudflare directly; the sidecar relays that traffic to a Shadowsocks server you run yourself abroad. Keeps Gluetun's kill-switch, healthchecks, and the entire rest of the stack (Tailscale, AdGuard, ipset, `routing-fix.sh`) untouched, since none of them know the transport changed underneath `warp`.

542

543 **\*\*B2 Replace:\*\*** Swap the `warp` service entirely for a Shadowsocks-client + `tun2socks` container that owns the network namespace instead of Gluetun. More invasive loses Gluetun's built-in kill-switch/healthcheck, which would need to be rebuilt by hand (see Section 12.9 for the pluggable architecture to support this natively).

544

545 **\*\*Recommendation: B1.\*\*** Smaller surface area, preserves the self-healing architecture already built and documented.

546

547 **### 12.6 Fingerprinting caveat**

548 Plain Shadowsocks can still be caught by active-probing DPI in more aggressive censorship environments. If plain SS gets blocked too, the next step up is an obfuscation plugin (`v2ray-plugin`, `Cloak`) or a newer protocol designed against active probing (VLESS+Reality, Hysteria2). Test plain SS first don't over-build before confirming it's needed.

549

550 **### 12.7 AmneziaWG (The High-Speed Alternative)**

551 If you want to maintain the bare-metal speeds of WireGuard while bypassing DPI (e.g. in Egypt or Russia), the best alternative to Shadowsocks/V2Ray is AmneziaWG. AmneziaWG is a fork of WireGuard that pads the handshake packets with randomized "junk" data, entirely destroying the 148-byte DPI signature that firewalls look for.

552 - **\*\*Pros:\*\*** Because it runs entirely as UDP without TCP-over-TCP overhead, it completely avoids "TCP meltdown" latency spikes associated with Shadowsocks/V2Ray wrappers.

553 - **\*\*Cons:\*\*** You cannot use Cloudflare WARP. Furthermore, it requires completely ripping Gluetun out of the `nullexit` stack and building a custom Docker container, meaning you lose Gluetun's built-in kill-switches and health-check logic.

554

555 **### 12.8 Infrastructure Costs & Privacy**

556 To use either Shadowsocks (Path B) or AmneziaWG, you cannot use the free Cloudflare WARP infrastructure, because Cloudflare servers only speak standard WireGuard. The software for both custom protocols is 100% free and open-source, but you must host the remote server infrastructure yourself.

557 - **\*\*Free Tier (Oracle Cloud):\*\*** You can host your remote server on Oracle Cloud's "Always Free" tier for \$0. However, this routes your highly sensitive, censorship-evading traffic directly through Oracle's massive corporate data broker. If privacy is the core objective, handing all your metadata to Oracle defeats the purpose.

558 - **\*\*Paid Tier (Hetzner / Mullvad):\*\*** The recommended path is to rent a standard ~\$4/month Virtual Private Server (VPS) in a free country (e.g., Hetzner in Germany, subject to strict EU GDPR privacy laws) and self-host the server. Alternatively, Mullvad VPN (5/month) provides native v2ray/Shadowsocks bridges built into their network, allowing you to bypass censorship with a strict zero-logs policy. It is highly recommended to pay with untrackable payment methods to maintain complete anonymity, which these providers typically support without requiring local residency.

559

560 **### 12.9 The Future: Egress Compartmentalization (Pluggable Architecture)**

561 To natively support invasive replacement paths like **\*\*Path B2\*\*** or **\*\*AmneziaWG\*\*** (Section 12.7) without breaking the core `nullexit` routing and monitoring logic, the architecture must be refactored into a modular, "pluggable" design.

562

```

563 1. The Egress Plugin System: Currently, WARP-specific logic (like `cdn-cgi/trace` health checks and
    hardcoded interface names) is scattered across `toggle.sh` and `scripts/routing-fix.sh`. This should be
    abstracted into a `scripts/plugins/` directory, where each egress method (`warp.sh`, `v2ray.sh`) implements
    standardized functions: `get_egress_interface()`, `check_health()`, and `get_bypass_ips()`.
564 2. Dynamic Docker Compose: Based on an `EGRESS_TYPE` variable in `.env`, dynamic compose file generation
    would spin up only the required egress containers (e.g., skipping the `warp` container when `EGRESS_TYPE=
    v2ray` is set).
565 3. Agnostic Routing & Watchers: `routing-fix.sh` would dynamically read the `EGRESS_INTERFACE` from the
    active plugin to apply `iptables` rules, and the background watchers in `toggle.sh` would become an
    agnostic `Egress Watcher` that invokes `check_health()` periodically.
566
567 This compartmentalization transforms `nullexit` into a universal, censorship-resistant gateway where the entire
    egress layer can be swapped by changing a single `.env` variable and running `./toggle.sh --restart`.
568
569 ---
570
571 ## 13. Roaming & Recovery Design: How nullexit Mimics Commercial VPNs
572
573 The `nullexit` roaming/network recovery subsystem replicates the exact engineering mechanics used by commercial
    , premium VPN clients (like Mullvad or NordVPN) to transition across Wi-Fi networks safely and seamlessly.
    This is the design overview; the specific bugs encountered and fixed along the way live in the Incident Log
    (15.5 Roaming/Sleep-Wake, 15.7 Tailscale P2P/SNAT).
574
575 ### The 4-Step Transition Cycle
576 When your Mac disconnects from one Wi-Fi and associates with another, `nullexit` coordinates the following
    events:
577
578 1. Firewall Lock (Kill-Switch): The macOS Packet Filter (PF) anchor `com.apple/nullexit` remains locked on
    the physical interface (`en0`), blocking all direct outbound IPv4 and IPv6 traffic. This guarantees that
    your real ISP IP or unencrypted DNS requests never leak onto the new network while the connection is
    rebuilding.
579 2. System Event Interception: A launchd LaunchAgent (`watcher.sh`) listens to the macOS System
    Configuration framework for the `State:/Network/Global/IPv4` key changes, spawning `recover.sh --post-wake`
    in under 500ms when a change is detected.
580 3. Bypass Routing: The script dynamically resolves the new physical network's IP gateway (e.g.
    `192.168.137.1`), cleans up stale bypass routes from the previous network, and adds static bypass routes
    pointing to Cloudflare's WARP endpoints and the Tailscale control plane directly via the new gateway. This
    allows the VPN client to negotiate its handshake outside of the tunnel.
581 4. Hole-Punching / NAT Re-negotiation: The script runs a target check on the `warp` container's tunnel. If
    the WireGuard connection is wedged by the network switch, it force-recreates the container to establish a
    fresh NAT binding to Cloudflare, restoring connection in ~15 seconds.
582
583 ---
584
585 ## 14. Deployment, Packaging & Known Limitations
586
587 ### 14.1 Packaging nullexit as a macOS Application (.dmg)
588
589 nullexit can be packaged into a native `.app` bundle, but this introduces nested virtualization requirements.
590
591 nullexit uses Colima (Apple Virtualization Framework `vz`) to run a Linux VM hosting Docker. Installing the
    `.app` inside a virtualized macOS environment (Parallels, cloud Mac) means running a Linux VM inside a
    macOS VM requiring hardware-level nested virtualization.
592
593 #### Hardware Support
594 - Supported: Apple Silicon M3/M4 (macOS Sequoia 15+), Intel Macs with VMX passthrough.
595 - Unsupported: M1, M2, A18 Pro. Colima crashes silently; the terminal (`setup.sh`) surfaces crash logs.
596
597 #### Build Steps
598 ```bash
599 # 1. Verify nested virtualization support
600 sysctl -a | grep hv_nested_virt_supported # expect: 1
601
602 # 2. Compile the AppleScript launcher
603 osacompile -o "Nullexit.app" "Toggle Gateway.applescript"
604
605 # 3. Embed scripts into app resources
606 mkdir -p "Nullexit.app/Contents/Resources/scripts"
607 cp toggle.sh setup.sh recover.sh docker-compose.yml "Nullexit.app/Contents/Resources/"
608
609 # 4. Package into DMG
610 hdiutil create -volname "Nullexit Installer" -srcfolder "./Nullexit.app" -ov -format UDZO Nullexit.dmg
611
612 # 5. Bypass Gatekeeper (unsigned app)
613 xattr -cr /Applications/Nullexit.app
614 ```
615
616 Without a paid Apple Developer certificate ($99/yr), users see an "App is damaged" Gatekeeper warning and need
    to run step 5.
617
618 ### 14.2 Moonlight/Sunshine + Tailscale Race Condition

```



```

619
620 When streaming from a remote Windows host running Sunshine over a Tailscale mesh (e.g., while the client is on
    a cellular hotspot), a race condition can occur on boot:
621
622 1. Both Tailscale and Sunshine are set to start automatically on boot.
623 2. Tailscale takes a few seconds to fully initialize its `100.x.x.x` virtual adapter and establish a connection
    .
624 3. **Sunshine starts too fast.** It launches before Tailscale is fully ready. Since Sunshine only binds to
    network interfaces that are active at the exact moment of startup, it misses the Tailscale adapter entirely
    .
625 4. Moonlight clients (even when configured with the correct Tailscale IP) will fail to connect because Sunshine
    isn't listening on that interface.
626
627 **The "Magic Toggle" Symptom:**
628 If the user manually enables or disables Wi-Fi on the host PC, it triggers a global Windows network change
    event. Sunshine listens for these events, re-enumerates all network adapters, and says, "Oh, there's a
    Tailscale adapter here!" and finally binds to it. Suddenly, Moonlight can connect.
629
630 **The Permanent Fix:**
631 On the Windows host, open `services.msc`, locate the **Sunshine Service**, and change its Startup type from **
    Automatic** to **Automatic (Delayed Start)**. This forces Windows to wait a minute or two after booting
    before launching Sunshine, ensuring Tailscale's virtual adapter is fully initialized first.
632
633 ### 14.3 Chrome Remote Desktop Performance (UDP NAT)
634
635 #### Observation (July 11, 2026)
636 When the user connects to their host Mac using Chrome Remote Desktop (CRD) while `nullexit` is active, the
    connection suffers from massive latency, lag, and degraded video quality. AdGuard Home's `blocked.log`
    shows zero blocked DNS requests for Google, WebRTC, or `chromoting` endpoints, indicating this is not a DNS
    sinkhole issue. *(See also 15.11 for the separate CRD-on-remote-device connection-failure quirk.)*
637
638 #### Root Cause
639 This is a fundamental limitation of tunneling all host traffic through a strict commercial NAT (Cloudflare WARP
    ).
640 1. Chrome Remote Desktop attempts to use **STUN (UDP hole-punching)** to establish a lightning-fast, direct
    peer-to-peer WebRTC connection between the client device and the host Mac.
641 2. Because all non-Tailscale traffic is forcefully routed through the `warp` container, the STUN packets exit
    the network from a Cloudflare IP.
642 3. Cloudflare's enterprise NAT infrastructure strictly drops unsolicited inbound UDP packets. The UDP hole-
    punch fails.
643 4. Because a direct P2P connection cannot be established, CRD falls back to **Relay Mode** (TURN), bouncing all
    video data back and forth through Google's cloud servers instead of sending it directly over the local or
    mesh network. This relaying introduces massive latency.
644
645 #### Workaround / Fix
646 This lag is the accepted "tax" for maintaining absolute cryptographic anonymity; the only way to restore CRD
    speed would be to leak its UDP traffic outside the tunnel (violating zero-trust).
647 To achieve low-latency remote access, users should bypass CRD entirely and use macOS's built-in **Screen
    Sharing (VNC)** directly over the Tailscale IP. Tailscale's sophisticated custom DERP/WireGuard
    architecture successfully UDP hole-punches through strict NATs, providing a direct, encrypted, high-speed
    connection without relying on external cloud relays.
648
649 ### 14.4 Dynamic IP Fetch vs. MagicDNS
650
651 #### Observation
652 The `nullexit` gateway uses a dynamically fetched IP address (cached in `.gateway_ip`) to configure host
    routing and the `pf` kill-switch, rather than utilizing Tailscale's user-friendly MagicDNS hostname (e.g.,
    `nullexit-gateway`).
653
654 #### Why MagicDNS is Not Used
655 1. **The "Chicken and Egg" Bootstrapping Problem:**
656 macOS's Packet Filter (`pf`) and low-level routing commands (`route add`) operate strictly at Layer 3 and
    require raw IP addresses. If `toggle.sh` or `recover.sh` attempted to use `nullexit-gateway`, the host Mac
    would need to perform a DNS lookup to resolve it. However, the very first step of the gateway's
    initialization is to completely hijack the host's DNS and point it **at the gateway itself**. The Mac cannot
    resolve the gateway's MagicDNS name if it needs the gateway's IP to know where to send the DNS query!
657 2. **Dynamic Self-Healing (Avoiding Stale Cache):**
658 To solve the bootstrapping problem without causing bugs if the node is deleted and re-authenticated on the
    Tailscale Admin Console, `toggle.sh` explicitly runs `docker compose exec ... tailscale ip -4` during boot
    to fetch the absolute freshest IP dynamically. It saves this to `.gateway_ip` so that fast-roaming scripts
    like `recover.sh` can instantly retrieve the routing target without incurring the latency of querying
    Docker.
659
660 ### 14.5 Battery & Power Measurement Quirks
661
662 When trying to optimize this framework (or any daemon) for battery life, developers often look for a way to
    programmatically measure the exact Wattage consumed by a specific process (e.g., "How many Watts is `toggle
    .sh` using?").
663
664 **This is a hardware impossibility on macOS and Linux.**
665

```



```

666 ##### Why Activity Monitor is Lying to You
667 Your machine's logic board only has physical power sensors for the *entire* CPU package, the GPU, and the RAM.
    It knows the CPU is pulling 5W total, but it has no physical hardware capability to know whether Docker is
    using 3W and Chrome is using 2W.
668
669 The "Energy Impact" score you see in macOS Activity Monitor is a **synthetic, heuristic score** calculated by `
    powerd`. It penalizes apps for waking up the CPU (idle wakes), doing disk I/O, or keeping the screen awake.
    It is a relative score, not actual Wattage.
670
671 ##### The Proper A/B Testing Methodology
672 If you want to truly know how many Watt-hours or mAh your background daemons (`routing-fix.sh`, the DNS/WARP
    watchers, etc.) are costing you, you must perform a hardware-level A/B drain test:
673
674 1. Unplug the laptop, run the gateway, and note your exact battery mAh using:
675     ```bash
676     ioreg -l | grep "AppleRawCurrentCapacity"
677     ```
678 2. Leave the laptop idle for exactly 1 hour.
679 3. Run the command again to see exactly how many mAh were drained.
680 4. Turn the gateway off and repeat the test for another hour.
681
682 The delta in mAh drained between the two tests is the true, exact cost of running the software. When optimizing
    polling loops (e.g., changing `sleep 5` to `sleep 30`), this is the only reliable way to measure the
    actual battery life returned to the user.
683
684 ### 14.6 Host Lockdown Mode (Why It's Impossible on macOS)
685
686 A common privacy feature request is a "Host Lockdown Mode" (Mode 3): A mode where the Docker exit node remains
    perfectly functional for remote devices, but the host machine itself is completely blocked from accessing
    the internet (to prevent even encrypted host traffic from leaking metadata).
687
688 **Why this is impossible on macOS:**
689 On Linux (e.g., a Raspberry Pi), Docker runs natively. The host's traffic traverses the `OUTPUT` iptables chain
    , while Docker traffic traverses the `FORWARD` chain. We could easily drop all `OUTPUT` traffic to kill the
    host's internet, and Docker would continue routing traffic perfectly.
690
691 However, Docker on macOS runs inside a Linux Virtual Machine (via `qemu` or Apple `Virtualization.framework`).
    This VM runs as a standard macOS user-space application. If you configure the macOS firewall (`pf`) to drop
    all host outbound traffic, it will indiscriminately drop the VM application's traffic as well, instantly
    killing the Cloudflare WARP tunnel and the exit node.
692
693 Writing complex macOS `pf` rules to "allow the VM app, allow Cloudflare WARP IPs, allow Tailscale DERP IPs, but
    block everything else" is highly fragile and heavily discouraged. Since the current `toggle.sh` default
    mode already fully encrypts the host's traffic via the WARP tunnel, the host is already protected from ISP
    leaks.
694
695 ---
696
697 # Part III Incident Log & Resolved Issues
698
699 This is the single, deduplicated log of every incident, bug, and resolved issue. Entries are organized into
    subsystem groups (15.115.12). Each entry follows a **Symptom / Root Cause / Fix** shape where applicable;
    packet-level analyses are preserved as clearly-labeled **Deep dive** blocks. The terse dated changelog
    lives in 17; this section holds the full post-mortems.
700
701 ## 15. Incident Log
702
703 ### 15.1 Exit-Node Return Path & SOCKS5 Failover
704
705 ##### 15.1.1 Exit Node Return Path (`rx 0`) & The SOCKS5 Failover RESOLVED
706 **Symptom:** Gateway showed `tx 936 rx 0` transmitted but never received return traffic. For days, Tailscale's
    exit node refused to return traffic back to the client, leading to the assumption that Tailscale's
    userspace forwarding was bypassing `iptables` and ignoring the WARP `tun0` interface. In a desperate
    attempt to fix this, a **SOCKS5 proxy** (`scripts/socks5-proxy.py`) was introduced as a replacement.
707
708 **Root Cause:** Tailscale's userspace forwarding was *not* the problem. `docker-compose.yml` included `
    FIREWALL_OUTBOUND_SUBNETS=100.64.0.0/10`. Gluetun created a strict routing rule (`ip rule add to
    100.64.0.0/10 lookup 199`, priority 99) that forced all Tailscale CGNAT traffic to bypass the VPN via the
    unencrypted Docker bridge (`eth0`). Return packets were sent to the macOS host instead of back through `
    tailscale0`, blackholing all return packets back to the client.
709
710 **Fix:** Removed `FIREWALL_OUTBOUND_SUBNETS` entirely, immediately restoring the Tailscale exit node. `scripts/
    routing-fix.sh` now injects `lookup 52`, and return packets correctly flow back into `tailscale0`. Instead
    of removing the SOCKS5 proxy, we repurposed it into a **bulletproof failover** for when restrictive Wi-Fi
    networks block Tailscale UDP ports.
711
712 Architecture after the fix:
713 ```
714 PRIMARY (Exit Node):
715 DNS/TCP/UDP: Apps  Tailscale Exit Node  gateway container  tun0  WARP  internet
716

```

```

717 FAILOVER (SOCKS5 - if Tailscale is blocked by local Wi-Fi):
718 DNS: Browser 127.0.0.1:53 Python proxy TCP:5354 AdGuard WARP
719 TCP: Apps SOCKS5:1080 gateway container tun0 WARP internet
720 Mesh: SSH/ping 100.x.x.x utun5 WireGuard peers (independent of proxy)
721 ```
722
723 **Deep dive: Exit Node Return Path Analysis (June 26, 2026).**
724 This preserves the detailed packet-level analysis that led to identifying and fixing the `rx 0` bug.
725
726 *The exit node was probably never going to work in this container topology.* Here's the packet path traced by
    reading `ip rule show` and the routing tables live:
727
728 **Outbound (Phone-1 internet):**
729 1. Phone-1 sends packet to gateway via Tailscale mesh
730 2. Gateway's tailscale0 (kernel TUN, `TS_USERSPACE=false`) decrypts packet enters FORWARD chain
731 3. Packet has src=`100.87.42.87` (Phone-1), dst=some internet IP
732 4. FORWARD chain (nftables): Rule 1 matches (`-i tailscale0 -j ACCEPT`)
733 5. Routing decision for the forwarded packet. Which ip rule matches?
734 - Rule 98: `to 172.18.0.0/16` no (dst is internet)
735 - Rule 99: `to 100.64.0.0/10` no (dst is internet)
736 - Rule 100: `from 172.18.0.2` **NO** (src is `100.87.42.87`, not the container IP!)
737 - Rule 101: `not fwmark 0xca6c lookup 51820` **YES** (no fwmark on forwarded packets)
738 6. Table 51820: `default dev tun0` packet goes out through WARP
739 7. NAT POSTROUTING: `MASQUERADE on tun0` src becomes WARP IP
740 8. Packet exits through WARP to internet (in theory)
741
742 **Return (internet Phone-1):**
743 1. Return packet arrives on tun0, conntrack de-MASQUERADES dst back to `100.87.42.87`
744 2. Routing decision for dst=`100.87.42.87`:
745 - Rule 99: `to 100.64.0.0/10 lookup 199` **YES**
746 3. Table 199: `100.64.0.0/10 via 172.18.0.1 dev eth0` **sends it to the Docker gateway!**
747 4. **This is wrong.** The packet should go to tailscale0 (to be re-encrypted and sent to Phone-1), but table
    199 sends it out eth0 to the Docker bridge host.
748
749 **Table 199 is the smoking gun.** It has one route: `100.64.0.0/10 via 172.18.0.1 dev eth0`. This sends ALL
    Tailscale CGNAT return traffic out the Docker bridge to the host not back through tailscale0. The return
    packet leaves the container's network namespace entirely, hits the host's routing stack, and gets dropped
    because the host has no idea what to do with a raw packet for `100.87.42.87`.
750
751 *Why table 199 exists and why it's wrong for exit node traffic.* Table 199 + the CGNAT rule (pref 99) was
    designed so that the **container itself** can reach other Tailscale peers (e.g., for mesh management, DERP
    relay). For that use case, routing CGNAT traffic via the Docker bridge to the host (which has its own
    tailscaled) makes sense. But for **forwarded exit node traffic**, the return packet needs to go back
    through tailscale0 inside the container not out to the host. The CGNAT rule intercepts the return packet
    before it can be routed back to tailscale0 through the main table or Tailscale's own routing (table 52,
    which is at pref 5270).
752
753 *The SOCKS5 path works because it dodges all of this.* SOCKS5 proxy creates connections **from the container's
    own IP** (`172.18.0.2`). So:
754 - ip rule 100 (`from 172.18.0.2 lookup 200`) matches
755 - Table 200 has `default dev tun0`
756 - Return traffic comes back to `172.18.0.2` on tun0, conntrack de-NATs, proxy sends response to client
757
758 No FORWARD chain involved. No CGNAT rule involved. No table 199. It just works.
759
760 *Summary of findings:*
761
762 | Finding | Status | Evidence |
763 |-----|-----|-----|
764 | Table 199 CGNAT route hijacks exit node return traffic | **Fixed** | Changed `lookup 199` to `lookup 52` in `
    scripts/routing-fix.sh`. Return packets now route via tailscale0. |
765 | FORWARD RELATED, ESTABLISHED missing | **Fixed** | Installed `iptables` via apk in `docker-compose.yml`. Rule
    now injects successfully. |
766 | DOCKER_GW variable parsing error | **Fixed** | Added `head -1` in `scripts/routing-fix.sh`. |
767 | `FIREWALL_OUTBOUND_SUBNETS` was the root cause | **Fixed** | Removed from `docker-compose.yml`. Gluetun no
    longer hijacks CGNAT routing. |
768 | Host tailscaled error state from aggressive `tailscale down` | **Mitigated** | Toggle script now detects and
    auto-restarts. |
769
770 #### 15.1.2 Container Default Route Bypasses WARP
771 **Problem:** Inside the WARP container's network namespace, the main routing table's default route was via `
    eth0` (Docker bridge), not `tun0` (WARP tunnel). While gluetun uses policy routing (rule 101 table 51820
    `tun0`) for most traffic, traffic originating from the container's own IP (`172.18.0.2`) hits rule 100 (`
    from 172.18.0.2 lookup 200`), whose default was also via `eth0`. This caused the SOCKS5 proxy's outbound
    connections to bypass WARP.
772
773 **Fix:** Updated the `routing-fix` container to:
774 1. Change the main table's default route to `default dev tun0`
775 2. Change table 200's default route to `default dev tun0`
776 3. Add a specific route for the WARP WireGuard endpoint (`162.159.192.1`) through `eth0` via the Docker gateway
    in table 200, preventing a tunnel loop

```

```

777 4. Re-assert all routes every 30 seconds (gluetun may reset them on health checks)
778 5. Dynamically detect the Docker subnet from `ip route show dev eth0` instead of hardcoding `172.18.0.0/16`
779
780 ##### 15.1.3 Hardcoded Docker Subnet in Routing Fix
781 **Problem:** The `routing-fix` container hardcoded `172.18.0.0/16` and `172.18.0.1` for the Docker bridge
       subnet and gateway. If Docker's bridge network changed (after `docker network prune`, Colima restart, or on
       a different machine), these routes would be wrong.
782
783 **Fix:** Detection is now dynamic using `ip route show`:
784 - `DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')
785 - `DOCKER_GW=$(ip route show default | awk '{print $3}')
786
787 This extracts the actual subnet and gateway directly from the routing table, working with any subnet size
       (`/16`, `/24`, `/20`, etc.).
788
789 ### 15.2 Routing Loops & Subnet Collisions
790
791 ##### 15.2.1 The Infinite Recursive Routing Loop (Hotspot Paradox)
792 **Problem:** A critical network topology crash occurs if the host Mac connects to a Mobile Hotspot (e.g., a
       Windows PC or Phone) that is also connected to the Tailscale mesh and has "Use Exit Node" enabled.
       Because the Mac is hosting the Exit Node, it routes its internet traffic to the Hotspot. The Hotspot
       intercepts the Mac's traffic on its way to the cellular tower and routes it back to the Exit Node (the
       Mac) via Tailscale. The Mac receives it, tries to send it out again, and the Hotspot intercepts it again.
       This creates an infinite, recursive encryption loop that instantly destroys network bandwidth and drops all
       connections.
793
794 **Fix:** This is a physical routing paradox that cannot be resolved in software. The upstream physical router
       providing internet to the Mac must be allowed to route its traffic directly to the physical WAN
       interface. The user must manually disable "Use Exit Node" on the upstream device providing the hotspot.
795
796 **Advanced Workaround:** If the upstream hotspot is a Windows machine and the user explicitly wants it to
       remain connected to the Exit Node, a permanent static route can be injected into the Windows routing kernel
       to forcefully bypass the Tailscale WFP interception for WARP packets. By binding the route to the physical
       adapter's Interface Index (IF), it becomes a permanent, roaming-aware bypass.
797
798 On Windows (Admin Command Prompt):
799 ```cmd
800 route print (find Interface List, note the IF number for the active Wi-Fi adapter, e.g. 15)
801 route -p add 162.159.192.1 mask 255.255.255.255 0.0.0.0 IF 15
802 route -p add 162.159.193.1 mask 255.255.255.255 0.0.0.0 IF 15
803 ```
804
805 On Linux (Root Terminal):
806 ```bash
807 ip route add 162.159.192.1 dev wlan0
808 ip route add 162.159.193.1 dev wlan0
809 ```
810
811 On macOS (Root Terminal):
812 ```bash
813 route add -host 162.159.192.1 -interface en0
814 route add -host 162.159.193.1 -interface en0
815 ```
816
817 *(Note: If the upstream router is an iPhone or Android device, you cannot inject static routes without
       jailbreak/root. You must simply disable "Use Exit Node" in the mobile Tailscale app).*
818
819 This punches a tiny hole straight through the paradox, letting the Mac's WARP packets escape to the physical
       internet while keeping the rest of the upstream router's traffic securely trapped in the Tailscale Exit
       Node.
820
821 ##### 15.2.2 Infinite Egress Routing Loops & Subnet Takeover Paradoxes
822 **Problem:** Two distinct network loops can break host egress connectivity entirely when the exit node is
       active:
823
824 1. **The Recursive Host-VM Tunnel Loop:**
825 - **Mechanism:** The host Mac's default route points to the Tailscale exit node (`utun`). The exit node
       container sends its encrypted tunnel packets (destined for the WARP endpoint `162.159.192.1`) back to the
       host via the VM NAT. Without a static route exception on the host Mac, the host Mac routes those packets
       right back into the exit node (`utun`), creating an infinite loop that blackouts all host network egress.
826 - **ARP-Scoping Pitfall:** Simply binding the bypass route to the interface (`route add -host 162.159.192.1
       -interface en0`) fails when the default route is deleted or redirected to `utun`. Since `162.159.192.1` is
       not local to the physical link, macOS fails to resolve it via ARP, dropping all tunnel packets.
827 - **Fix:** Dynamically resolve the active physical gateway IP (e.g. `172.17.0.1`) on startup, and add the
       static host routes for the WARP endpoints via the gateway IP directly (`route add -host 162.159.192.1 <
       gateway_ip>`). Additionally, because standalone `tailscaled` on macOS does not automatically modify the
       system default route via the CLI, we manually redirect the default gateway to `utun` using `
       setup_exit_node_routing`.
828
829 2. **The Docker Compose Subnet Takeover Collision:**

```

```

830 - **Mechanism:** To avoid collisions with the host's campus network, Colima's default bridge (`docker.bip`)
      is moved (e.g. to `10.200.0.1/24`). However, because `172.17.0.0/16` is now free inside the VM, Docker
      Compose's IPAM dynamically takes it over for the project's custom network (`nullexit_default`). Because the
      containers receive IPs on `172.17.0.0/16`, the host Mac cannot send local peer discovery packets (e.g.,
      Tailscale `magicsock` packets on `172.17.0.2:41641`) directly. The host routing table sends them out `en0`
      to the physical Wi-Fi, returning `host is down` or `no route to host`.
831 - **Fix:** Explicitly configure a custom default network subnet `10.200.1.0/24` in `docker-compose.yml` to
      prevent Docker Compose from ever reclaiming the conflicting `172.17/172.18` ranges.
832
833 > See also 15.2.1 (The Hotspot Paradox) for the related infinite loop that occurs when the physical upstream
      router is also a Tailscale exit-node client.
834
835 ##### 15.2.3 Docker Default Bridge vs Host Wi-Fi Subnet Collision (The 172.17.0.0/16 Subnet Overlap)
836 **Context:** When attempting to ping a local Windows client (`172.17.22.187`) or the default gateway
      (`172.17.0.1`) directly over Wi-Fi, the Mac host returns `No route to host` (displaying a `REJECT` / `!`
      flag in the routing table). Tailscale on the client also drops offline shortly after starting when using
      the exit node, failing to establish direct P2P connections and falling back to slow DERP relays.
837
838 **Root Cause & Subnet Overlap Mechanism:**
839 1. **Docker Default Subnet:** By default, the Docker daemon initializes its default bridge network (`docker0`)
      on the `172.17.0.0/16` IP range.
840 2. **Subnet Conflict:** The host's physical Wi-Fi network happens to be assigned the exact same
      `172.17.0.0/16` range by the local network router.
841 3. **Routing Clash:** Colima launches a Linux VM on the Mac host to run the Docker engine. Because Docker
      inside that VM creates `docker0` and claims the `172.17.0.0/16` subnet, any packet sent to `172.17.x.x`
      from within the containers (such as Tailscale local discovery) is captured by the VM's internal virtual
      bridge interface (which is down/empty) and is blackholed. On the Mac host, macOS dynamically marks routes
      as `REJECT` (!) when local ARP resolution fails, causing local pings to immediately abort with `No route
      to host`.
842 4. **Tailscale Relay Saturation:** Because local peer-to-peer discovery is blackholed, Tailscale falls back to
      public DERP relay servers (e.g. `relay "tor"` in Toronto). When exit-node routing is enabled, all client
      web traffic is routed through the relay. Public relays impose strict rate limits and throttle high-volume
      traffic, resulting in packet drops. This saturation drops Tailscale's control packets, causing the
      connection to time out and showing the client as offline.
843
844 **Fix:**
845 The canonical script for this is `scripts/fix-docker-bridge-collision.sh`. **One command applies the entire fix
      :**
846
847 ```bash
848 bash scripts/fix-docker-bridge-collision.sh          # apply
849 bash scripts/fix-docker-bridge-collision.sh --dry    # preview changes without applying
850 ```
851
852 It auto-detects the host's actual Wi-Fi subnet (no more guessing whether the campus DHCP handed out `172.x`,
      `10.x`, or `192.168.x`), picks a non-conflicting `docker.bip` from a small candidate set (defaulting to
      `172.26.0.1/24` if no overlap is matched), backs up `~/colima/default/colima.yaml` with an ISO-8601
      timestamp, edits the file **in place** (idempotent replaces an existing `docker.bip` if present, appends
      under an existing `docker:` block, or creates a fresh block), restarts Colima, rebinds Wi-Fi DHCP, verifies
      the new default route is no longer Docker's bridge, and re-runs both `./toggle.sh` and `scripts/diagnose-
      host-leak.sh` to print the post-fix verdict.
853
854 For reference, the underlying manual fix is: edit the Colima configuration file (`~/colima/default/colima.yaml`
      ) on your Mac to define:
855 ```yaml
856 docker:
857   bip: 172.26.0.1/24    # any /24 that does NOT overlap your Wi-Fi subnet
858   ...
859 Then, restart Colima to apply the new subnet inside the VM:
860 ```bash
861 colima restart
862 ```
863 This forces Docker to use `172.26.x.x` for its default bridge, freeing the physical `172.17.x.x` range and
      letting pings and direct Tailscale peer-to-peer connections flow normally.
864
865 ##### 15.2.4 July 4, 2026: The "Network Unreachable" Host Leak & Post-Wake Container Destructions
866 **Symptom:**
867 1. Running `docker compose config` with dummy keys corrupted the user's `.env` file by appending invalid
      WireGuard keys. This caused the `warp` container to become unhealthy on boot, which triggered a full
      teardown of the gateway by `toggle.sh`.
868 2. After fixing `.env`, `toggle.sh` successfully booted the gateway, but the host's internet started bypassing
      the tunnel entirely (a massive host leak). The Cloudflare trace returned `warp=off` and exposed the host's
      physical ISP IP.
869 3. Running `recover.sh --post-wake` would violently force-recreate all gateway containers and drop the host's
      internet connection.
870
871 **Root Cause:**
872 - **Bug 1 (The Host Leak):** In `toggle.sh`, the logic to find the Tailscale `utun` interface used `ifconfig |
      grep -B4 "inet 100."`. Due to the 4 lines of before-context, this regex was accidentally capturing the
      interface **above** the actual Tailscale interface (e.g. grabbing `utun4` instead of `utun5`). Because `utun4`
      had no IP address, the subsequent `sudo route add -net 0.0.0.0/1 -interface utun4` silently failed with "

```

```

Network is unreachable", leaving the host's default route exposed.
873 - **Bug 2 (The Post-Wake Destructions):** The health check in `recover.sh --post-wake` relied on `docker
compose exec -T warp curl -sf https://www.cloudflare.com/cdn-cgi/trace`. However, the newer `qmcgaw/glutun
` Alpine-based Docker image no longer ships with `curl` pre-installed. The health check failed with `
executable file not found in $PATH`, tricking `recover.sh` into thinking the tunnel was completely dead and
needed to be forcefully recreated via `docker compose up -d --force-recreate`.
874 - **Bug 3 (Restart Flags):** Previous refactors incorrectly permitted `toggle.sh restart` instead of strictly
enforcing `toggle.sh --restart`.
875
876 **Resolution:**
877 - Cleaned the `.env` file of duplicate dummy entries.
878 - Replaced the brittle `grep -B4` interface parsing in `toggle.sh` with a strict AWK parser: `awk '/^[a-z0
-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':'`. This mathematically isolates the exact
interface possessing the Tailscale `100.x.x.x` IP address, preventing the `0.0.0.0/1` route from silently
failing.
879 - Changed the health check in `recover.sh` to use `wget -q0- --timeout=3` which is natively installed in the
Glutun image. This stopped the unnecessary destructive recreations of the containers during post-wake
events.
880 - Reverted the `restart` alias in `toggle.sh` to strictly require `--restart`.
881
882 #### 15.2.5 July 4-5, 2026: Tailscale Data-Plane Loop (The DERP Relay Deadlock)
883 **Symptom:**
884 After bypassing the Tailscale control plane (`192.200.0.0/16`) to fix the "offline" mesh status, the Mac
appeared online. However, external peer devices (like the user's phone on cellular) could not successfully
ping or SSH into the Mac. `tailscale netcheck` reported `Nearest DERP: unknown (no response to latency
probes)`.
885
886 **Root Cause:**
887 While bypassing the control plane kept the daemon connected to the coordination servers, Tailscale's actual
peer-to-peer data plane relies on STUN (for NAT traversal) and DERP relays. These relays span ~80
dynamically allocated IP addresses across multiple cloud providers. Because we were forcing `0.0.0.0/1`
through the `utun*` exit node, the Mac's own outbound connection attempts to DERP relays were being routed
back into the tunnel. To prevent an infinite loop, Tailscale's daemon dropped its own packets when it saw
them returning on the TUN interface. This effectively left the daemon deaf and blind to peer connections.
888
889 **Resolution:**
890 - Modified `add_warp_bypass_routes` in `scripts/common.sh` to execute a live API fetch (`curl -s https://login.
tailscale.com/derpmap/default`) during startup.
891 - The script parses out all active DERP relay IPv4 addresses (~80 IPs) and adds a physical host bypass route
for every single one of them.
892 - These IPs are written to a temporary file (`/tmp/nullexit-derp-ips.txt`) so they can be cleanly un-routed by
`remove_warp_bypass_routes` when the gateway shuts down. Peer connectivity (ping, SSH, SFTP) was fully
restored.
893
894 #### 15.2.6 July 5, 2026: Hardcoded WARP Endpoints in docker-compose
895 **Symptom:**
896 The bash scripts allowed overriding the default Cloudflare WARP IP endpoints via `.env` variables (`
WARP_ENDPOINT_1` and `WARP_ENDPOINT_2`) to establish bypass routes. However, `docker-compose.yml`
statically hardcoded `162.159.192.1` for the `warp` container's `WIREGUARD_ENDPOINT_IP`.
897
898 **Root Cause & Risk:**
899 If a user set a custom endpoint in `.env`, the script would successfully bypass the custom IP on the host level
. However, Glutun would still attempt to connect to the hardcoded `162.159.192.1`. Since `162.159.192.1`
was no longer explicitly bypassed, its packets would be sucked into the `utun*` exit node, causing an
infinite WireGuard routing loop that instantly breaks the gateway.
900
901 **Resolution:**
902 Modified `docker-compose.yml` to dynamically read the environment variable via `${WARP_ENDPOINT_1
:-162.159.192.1}`, ensuring both the host routing scripts and the container use the exact same endpoint.
903
904 #### 15.2.7 Incident Post-Mortem: WARP Watcher Race vs Post-Wake (July 10, 2026)
905 **Symptom:** After switching Wi-Fi networks, the host IP appeared as the raw ISP IP (`132.x.x.x`) instead of a
Cloudflare/WARP IP. The gateway appeared to complete post-wake recovery successfully in the logs, but
nuclear shutdown fired immediately after and killed everything.
906
907 **Root Cause:** A fatal race condition between two concurrent processes:
908
909 1. **Wi-Fi roam** `watcher.sh` fires `recover.sh --post-wake`
910 2. **Post-wake** finds the `warp` container missing/dead calls `docker compose up -d --force-recreate` (~35
seconds)
911 3. **In parallel**, the WARP Watcher (running as a child of `toggle.sh`) polls the warp container every 5s
during failures
912 4. While the container is being recreated it returns `warp=off` for those 35 seconds 7 consecutive failures
**WARP SHUTDOWN fires**, calling nuclear `recover.sh`
913 5. Nuclear `recover.sh` tears down Tailscale, resets DNS to 1.1.1.1, stops all containers **gateway dead, IP
exposed**
914 6. The post-wake `docker compose up` finishes at ~35s mark, container healthy but the gateway is already gone
915
916 The evidence in `output.log`:
917 ...
918 [2026-07-10T06:07:27Z] NET: ... bash recover.sh --post-wake

```

```

919 [2026-07-10T06:07:47Z] WARP DOWN   cdn-cgi/trace reports warp=off   watcher starts counting
920 ...
921 add net 0.0.0.0: gateway utun5     post-wake successfully re-routes at 02:08:22
922 ...
923 [2026-07-10T06:09:04Z] WARP SHUTDOWN 6 consecutive failures     watcher fires nuclear
924 ...
925
926 **Fix (July 10, 2026):** Added a WARP Watcher inhibit marker (`/tmp/nullexit-warp-inhibit.marker`):
927 - `recover.sh --post-wake` writes the marker at startup before any container operations
928 - The WARP Watcher loop checks for the marker at the top of every iteration; if present it resets `consec_off`
929   = 0 and sleeps 5s (skips the poll entirely)
929 - `recover.sh` removes the marker at the very end (on both success and failure via `trap cleanup_inhibit EXIT`)
930 - This guarantees the watcher never counts failures during the intentional container-down window of a post-wake
   force-recreate
931
932 The marker file path is hardcoded to `/tmp/nullexit-warp-inhibit.marker` in both `toggle.sh` and `recover.sh`.
933
934 #### 15.2.8 Incident Post-Mortem: Concurrency Collision between Nuclear Teardown and Post-Wake (July 10, 2026)
935 **Symptom:** When the background WARP Watcher triggers a nuclear shutdown (due to a real or simulated tunnel
   failure), the gateway is completely torn down. However, the network configuration changes made during the
   teardown trigger a `State:/Network/Global/IPv4` network change event. The LaunchAgent-based network watcher
   (`watcher.sh`) intercepts this event and concurrently spawns `recover.sh --post-wake` to heal/restore the
   gateway. This causes `docker compose down` (from nuclear teardown) and `docker compose up` (from post-wake)
   to run in parallel, resulting in container errors like `dependency failed to start: container warp exited
   (0)` and the gateway becoming wedged.
936
937 **Root Cause:** A lack of coordination/locking between the lifecycle control scripts. While `toggle.sh` used `/
   tmp/nullexit-toggle.lock` to prevent concurrent `toggle.sh` runs, `recover.sh` (both nuclear and `--post-
   wake` modes) did not utilize or check any lock files.
938
939 **Fix (July 10, 2026):** Implemented a shared lifecycle lock file (`/tmp/nullexit-toggle.lock`) across both
   scripts:
940 - **`recover.sh` (all modes):** Checks `/tmp/nullexit-toggle.lock` at startup. If the lock is held by another
   running lifecycle script (`toggle.sh` or `recover.sh`):
941   - In `--post-wake` mode, it logs a skip message to `output.log` and exits cleanly (`exit 0`). This safely
   prevents the auto-recovery agent from recreating containers during a teardown.
942   - In nuclear recovery mode, it exits with an error.
943 - **`recover.sh` Lock Cleanup:** Cleans up the lock file upon completion using `trap cleanup_recover EXIT` (
   which checks if the lock file belongs to the current PID).
944 - **`toggle.sh` update:** The lock check in `toggle.sh` was updated to look for both `toggle.sh` and `recover.
   sh` in the running processes to enforce proper exclusion.
945
946 #### 15.2.9 Incident Post-Mortem: Infinite Auto-Recovery Loop after WARP Watcher Nuclear Shutdown (July 10,
   2026)
947 **Symptom:** When the background WARP Watcher triggered a nuclear shutdown (due to a tunnel outage), it
   successfully executed `recover.sh` (nuclear). However, immediately after the teardown completed, the system
   started spawning `recover.sh --post-wake` and rebuilding the containers again. This created an infinite
   loop of tearing down and auto-restarting the gateway, keeping the host's internet route permanently wedged.
948
949 **Root Cause:** An active-state marker file leak. The LaunchAgent-based network watcher (`watcher.sh`) only
   fires `recover.sh --post-wake` when `/tmp/nullexit-gateway-active.marker` is present on disk. While `toggle
   .sh` (STOP path) correctly deleted this marker, the nuclear `recover.sh` script did not. When the WARP
   Watcher ran `recover.sh` (nuclear), the containers were stopped but the active-state marker was left on
   disk. The network changes from the teardown then triggered `watcher.sh`, which saw the marker, believed the
   gateway was still supposed to be active, and triggered `--post-wake` to start it back up.
950
951 **Fix (July 10, 2026):** Modified the start of `recover.sh` to explicitly delete `/tmp/nullexit-gateway-active.
   marker` if `POST_WAKE` is `false` (nuclear mode). This ensures that any subsequent network configuration
   changes made during the teardown are correctly ignored by `watcher.sh` because the marker file is gone.
952
953 ### 15.3 MTU, Fragmentation & Throughput
954
955 #### 15.3.1 Network Bufferbloat & MTU Optimizations (Double-VPN UDP Crash)
956 **Symptom:** When running aggressive UDP bandwidth tests (such as macOS `networkQuality`, which uses QUIC/HTTP3
   ), the `nullexit` tunnel completely crashes or exhibits severe lag (6,000ms+ bufferbloat).
957
958 **Root Cause:** The bug is a cascading MTU mismatch. Tailscale generates a UDP packet. Cloudflare WARP (also
   WireGuard) receives the packet, wraps it in another layer of encryption, pushing the packet size beyond the
   standard 1500 byte limit. Unlike TCP, UDP packets cannot be resized dynamically via MSS negotiation,
   resulting in massive IP packet fragmentation. The Apple Virtualization framework (`vz`), which Colima uses
   for bridged networking, silently drops heavily fragmented UDP packets under high load. This blackholes the
   entire UDP upload stream, instantly dropping the connection.
959
960 **Fix (Partial):** To prevent MTU fragmentation for standard web traffic, TCP MSS Clamping is strictly
   enforced via the macOS `pf` firewall. By adding `scrub out all max-mss 1160` to `scripts/pf.conf`, macOS
   unilaterally shrinks all outgoing TCP headers to fit comfortably inside the double-encrypted tunnel,
   bypassing Path MTU Discovery blackholes. `TS_DEBUG_MTU=1200` was also added to `docker-compose.yml` to
   minimize internal fragmentation.
961 **Unresolved:** The only way to solve the UDP crash bug natively is to artificially cap the maximum tunnel
   bandwidth using SQM (`fq_codel`). Because a static bandwidth cap would artificially throttle high-speed
   networks, `nullexit` leaves the bandwidth uncapped, optimizing perfectly for TCP while accepting UDP

```



fragmentation under extreme edge-case load tests.

962 **\*Why not Dynamic SQM (Aurate)?\*** Implementing an EMA/PID-controlled aurate daemon (like `sqm-aurate`) requires a constant 100ms polling loop to measure the kernel packet queue, which severely degrades laptop battery life by preventing deep C-state CPU idling. Furthermore, macOS's native firewall (`dummynet`) only supports `fq\_codel` and does not support the newer Linux `CAKE` algorithm, which is the only algorithm that offers kernel-native, event-driven `aurate-ingress` (zero polling penalty). Therefore, dynamic SQM is computationally hostile to battery-powered macOS hosts.

963

964 **#### 15.3.2 Double-Tunneling MSS Clamping (Stalling Bug)**

965 **\*\*Problem:\*\*** Traffic double-wrapped through Tailscale (WireGuard) and WARP (WireGuard) suffered 120 bytes of MTU overhead (60 + 60). The default MSS clamp of 1180 was still slightly too large, causing large packets to fragment or drop, which resulted in mysterious web page stalls on strict-MSS endpoints.

966

967 **\*\*Fix:\*\*** Lowered the TCP MSS clamp in `post-rules.txt` to `1120` to guarantee double-tunneled payloads fit safely within standard internet MTUs.

968

969 **#### 15.3.3 Wi-Fi Edge Packet Loss via QUIC (UDP) Fragmentation**

970 **\*\*Problem:\*\*** Applications that heavily utilize HTTP/3 (QUIC) over UDP such as Facebook, Instagram, and Google services experience massive stuttering and stalled image loading when the client device is far away from the router (weak Wi-Fi signal). The issue did not occur when close to the router.

971

972 **\*\*Root Cause:\*\*** The `nullexit` gateway uses a double-encrypted tunnel (Tailscale + WARP). To prevent packets from fragmenting when entering these tunnels, a firewall rule in `post-rules.txt` uses `--set-mss 1120` to safely clamp TCP packets. However, because QUIC runs entirely over UDP, it bypasses the TCP MSS handshake completely. QUIC blasts maximum-MTU (1280 byte) UDP packets. The inner `tailscale0` interface (MTU 1200) forces these large UDP packets to fragment into two pieces. When the client is on the edge of Wi-Fi range (where 5-10% packet loss is common), losing either UDP fragment destroys the entire image frame. This exponential amplification of packet loss stalled QUIC streams.

973

974 **\*\*Fix & Trade-offs:\*\*** Appended an `iptables` rule to `post-rules.txt` that explicitly blocks `UDP --dport 443` (QUIC). When modern web apps detect QUIC is blocked, they instantly fall back to HTTP/2 (TCP 443). Because TCP is routed through the MSS clamp, the packets are resized before they leave, entirely eliminating IP fragmentation and restoring high-speed loading over weak Wi-Fi.

975 **\*Consequences:** This adds a minor latency penalty (TCP 3-way handshake) compared to QUIC's zero-RTT connection. It may also force WebRTC video calls (Google Meet/Discord) to fall back to TCP TURN servers, slightly inflating real-time video latency. However, in a constrained double-tunnel mesh, the absolute stability gained from eliminating fragmentation vastly outweighs these trade-offs.

976

977 **#### 15.3.4 Cellular Networks & PMTUD Blackholes (The 1280 Byte Limit)**

978 **\*\*Experiment:\*\*** Conducted a live packet sweep from the PC-1 host to Phone-1 connected via a cellular network + Tailscale DERP relay using `ping -D -s <size>`.

979

980 **\*\*Result:\*\***

981 - Packets with a payload of `1252` bytes (+ 28 bytes IP/ICMP headers = `1280` bytes total) succeeded perfectly with ~60ms latency.

982 - Packets with a payload of `1253` bytes (`1281` bytes total) and above resulted in a silent timeout.

983

984 **\*\*Analysis:\*\*** The cellular network (or the Tailscale DERP relay) enforces a strict MTU of 1280 bytes. Critically, it acts as a "PMTUD Blackhole" it silently drops packets larger than 1280 bytes instead of returning an ICMP "Fragmentation Needed" warning. If the exit node tries to send a standard 1500-byte internet packet to the phone over this link, it vanishes, causing web pages to stall permanently.

985

986 **\*\*Conclusion:\*\*** This empirically validates the absolute necessity of the TCP MSS clamping rule in `post-rules.txt`. By artificially clamping the MSS to `1120` (or `1180`), we ensure the TCP payload + headers + double-WireGuard overhead never exceeds the strict 1280-byte ceiling of the cellular mesh link.

987

988 **#### 15.3.5 Low Gateway Throughput (DERP Relay & MTU Fragmentation) Fixed**

989 **\*\*Problem:\*\*** The host Mac's internet throughput through the gateway dropped significantly (e.g., from 34Mbps raw to 2.1Mbps via gateway). This was caused by a combination of two bottlenecks:

990

991 **\*\*1. Tailscale DERP Relay Bottleneck:\*\***

992 Because the host Mac and the Docker container operate behind the same public IP, standard hole-punching for Tailscale P2P failed, forcing traffic through Tailscale's NYC DERP relay. Normally, `routing-fix.sh` explicitly drops container Tailscale UDP packets destined for local subnets (when `TAILSCALE\_ALLOW\_LAN\_P2P=false` in `.env` or auto-detected). We nuke these P2P connections on purpose to prevent the severe macOS `gvproxy` SNAT endpoint poisoning issue we faced earlier (see 15.7 SNAT Endpoint Poisoning). However, this strict policy unintentionally blocked direct communication between the container and the Mac host itself via the Docker bridge network.

993

994 **\*\*Fix:\*\*** Tailscale's control plane registers the Mac host using its physical IP (e.g., `172.17.52.223`), not the Docker bridge IP. If this physical IP falls within the dropped local subnets (like `172.16.0.0/12`), the P2P handshake is dropped. To solve this natively, `scripts/common.sh` now features a `write\_host\_ips()` function that actively grabs all physical IPv4 addresses on the Mac and writes them to `.host\_ips`. `routing-fix.sh` dynamically reads this file every 30 seconds and injects priority `RETURN` rules for those exact physical IPs. This guarantees direct container-to-host P2P over the local bridge (bypassing DERP) while continuing to block external LAN devices (like phones) to prevent the 15.7 SNAT poisoning bug.

995

996 **\*\*2. MTU Double-Encapsulation Fragmentation:\*\***

997 The host's Tailscale interface (`utun5`) defaulted to an MTU of 1280. The WARP tunnel inside the container also uses an MTU of 1280. When a full 1280-byte Tailscale packet from the host entered the container and was encapsulated by WARP (adding ~60 bytes of overhead), the resulting packet exceeded 1280 bytes, leading to



```

severe fragmentation and performance degradation.
998
999 **Fix:** In `toggle.sh`'s `setup_exit_node_routing` function, the host's Tailscale `utun` interface MTU is
explicitly lowered to `1200` (configurable via `HOST_MTU` in `.env`). This ensures that host-originated
packets can comfortably fit inside the container's WARP tunnel encapsulation without fragmenting.
1000
1001 > **Design Note Is direct hostcontainer P2P a Docker bug, or are we "exploiting a VM/host relationship"?
1002 > Neither. It is expected, well-defined behavior, and the RETURN-rule carve-out is a *defensive de-restriction
*, not an exploit. The reasoning:
1003 > - **On native Linux**, the container and host share one kernel; the Docker bridge is a first-class local
interface and hostcontainer is *trivially* direct. Nobody calls that an exploit. Our RETURN rules simply
reproduce that same local reachability on macOS.
1004 > - **The RETURN rules add nothing new; they *stop blocking* something legitimate.** `routing-fix.sh` broadly
DROPS containerLAN Tailscale UDP to prevent the real 15.7 `gvproxy` SNAT endpoint-poisoning bug caused by *
external* LAN peers (e.g. a phone on another AP). Tailscale's control plane registers the Mac host by its *
physical* LAN IP (e.g. `172.17.52.223`), which falls inside those DROPPed ranges so the host, the one peer
that is literally the same physical machine, became collateral damage. The `.host_ips` RETURN rules are
the minimal, precise exception that re-permits exactly that same-machine path and nothing else.
1005 > - **The genuinely VM-specific wart is `gvproxy` SNAT poisoning (15.7), and we route *around* it, not through
a hole in it.** Keeping hostcontainer traffic on the Docker bridge (direct) avoids letting it get NATed
through the userspace `gvproxy` endpoint rewriter that confuses Tailscale's peer-endpoint learning. That is
the opposite of exploiting the VM boundary it is respecting it.
1006 > - **What genuinely *doesn't* work on a VM host is UDP hole-punching to arbitrary *external* peers** (see
15.13). That is a real limitation of userspace VM networking, and it is why remote peers fall back to DERP
unless bridged mode is enabled on a trusted LAN. Direct hostcontainer P2P is the one case that works
regardless because both endpoints live on the same physical machine, no hole-punching is needed at all. It
was verified `direct 172.17.52.223` even in plain user-mode networking (no `--network-address`).
1007 >
1008 > In short: the DROP is the deliberate policy, the RETURN is a surgical exception for same-machine traffic, and
the whole thing is standard bridge networking not a Docker bug and not an exploit.
1009
1010 ### 15.4 DNS Interception & Proxy
1011
1012 #### 15.4.1 DNS Proxy: TCP Wire Format Mismatch (socat / dnsmasq)
1013 **Problem:** When the Tailscale data plane is unavailable (DERP relay only), the script falls back to a local
DNS proxy that forwards queries from host UDP:53 to Docker's TCP:5354 (AdGuard). The `socat` and `dnsmasq`
approaches both failed.
1014
1015 - **socat** forwards raw bytes it does not prepend the 2-byte length prefix that DNS-over-TCP requires.
AdGuard parses the stream incorrectly and returns garbage.
1016 - **dnsmasq** tries UDP first to reach the upstream (`127.0.0.1#5354`), but Colima's SSH tunnel only forwards
TCP. With a single upstream server, dnsmasq doesn't fall back to TCP even with `--timeout=1`.
1017
1018 **Solution:** Replaced both with a **Python DNS proxy** (`scripts/dns-proxy.py`, ~25 lines) that properly
handles the DNS-over-TCP wire format: reads a UDP query from the host, prepends the 2-byte length prefix,
sends it over TCP to AdGuard, strips the prefix from the response, and sends it back over UDP.
1019
1020 *Architectural Note: Why is the DNS proxy in Python while the SOCKS5 proxy was moved to a compiled Go Docker
container?
1021 The SOCKS5 proxy handles heavy data streaming (e.g., video, downloads). Python introduces massive CPU/battery
overhead for raw socket streaming, which is why we rely on the compiled Go implementation built directly
into the `tailscaled` daemon via `TS_SOCKS5_SERVER=:1080`.
1022 Conversely, the DNS proxy only handles intermittent UDP queries (a few bytes per minute) generated solely by
the local host machine. The Python overhead for this is effectively 0.001 Watts. We deliberately kept `dns-
proxy.py` in Python because it runs natively on the macOS/Linux hostrewriting it in Go would force users to
install a Go compiler (`brew install go`) on their host machine for zero noticeable battery gain. (Note:
If this architecture is ever adapted to serve as a public DNS resolver for thousands of users or for
massive automated web-scraping clusters, `dns-proxy.py` must be rewritten in Go to survive the concurrent
packet flood).
1023
1024 #### 15.4.2 Tailscale MagicDNS Bypassing the AdGuard PREROUTING Intercept
1025 **Context:** The `routing-fix.sh` script applies an iptables NAT rule (`PREROUTING -i tailscale0 -p udp --dport
53 -j REDIRECT`) to intercept all incoming DNS queries from connected Tailscale peers and force them into
the AdGuard container on port 5335.
1026 **Problem:** When remote clients (like Android or iOS) have "Use Tailscale DNS settings" turned on, they query
Tailscale's MagicDNS IP (`100.100.100.100`). This packet enters the exit node, but the `tailscaled` daemon
running *inside* the gateway container intercepts the `100.100.100.100` packet internally. `tailscaled`
then generates a brand new, local DNS request to the upstream DNS server to resolve the query. Because this
new request is generated *locally* by the daemon (not arriving externally via the `tailscale0` interface),
it bypasses the `PREROUTING` chain completely. The request glides straight out the WARP tunnel to
Cloudflare/Google, entirely bypassing the AdGuard sinkhole.
1027 **Fix:** The only way to fix this blindspot is to force `tailscaled` to use AdGuard as its upstream. In the
Tailscale Admin Console, the gateway's Tailscale IPv4 address (e.g., `100.x.x.x`) must be added as the sole
Custom Global Nameserver, and "Override local DNS" must be enabled. Additionally, to block Android devices
from bypassing this via Private DNS (DNS-over-TLS), outbound Port 853 traffic is explicitly REJECTED in
both `routing-fix.sh` and `post-rules.txt`.
1028
1029 #### 15.4.3 Seamless Wi-Fi Roaming & The macOS DNS Wipeout Loophole
1030 **Problem:** The Tailscale, WireGuard, and Colima NAT stacks are natively designed for connectionless roaming
and will seamlessly survive when the macOS host switches Wi-Fi networks (e.g., roaming from home Wi-Fi to a
cellular hotspot). However, the internet completely breaks upon network transition. This occurs because

```

macOS intentionally wipes out all custom DNS settings (`networksetup -setdnsservers`) and reverts to the new network's default DHCP nameserver whenever the active Wi-Fi BSSID changes. Since the Mac is locked to the exit node, its DNS requests are swallowed by WARP and dumped onto the public internet, which cannot route private DHCP IPs. This results in a total DNS failure.

1031

1032 **\*\*Solution:\*\*** Implementing a silent background "DNS Watcher" daemon (`nullexit-dns-watcher`) in `toggle.sh`. When the gateway starts, it spawns a background polling loop that checks the active Wi-Fi DNS every 30 seconds. If macOS resets the DNS during a network change, the watcher instantly detects the drift and forcefully re-injects `100.100.21.8` via `networksetup`. When the gateway stops, the watcher process is cleanly killed. This allows true, seamless roaming across physical networks without ever dropping the gateway state.

1033

1034 **#### 15.4.4 docker compose exec `r` Injection**

1035 **\*\*Fix:\*\*** All `docker compose exec` output piped through `tr -d 'r'`.

1036

1037 **#### 15.4.5 Stderr Invisibly Swallowed**

1038 **\*\*Symptom:\*\*** Failures were silent due to stderr redirected to `output.log`.

1039

1040 **\*\*Fix:\*\*** Tailscale wait loop now shows output. Critical commands show errors inline.

1041

1042 **## 15.5 Roaming, Sleep/Wake & Auto-Recovery**

1043

1044 **#### 15.5.1 Gateway Breakage on Lid Close / Wi-Fi Roam Post-Wake + Post-Roam Auto-Recovery**

1045 **\*\*Context / Symptom:\*\*** Closing the lid (sleep/wake) OR losing + reconnecting Wi-Fi (e.g. in an elevator, caf, or roaming between APs) leaves the gateway in a partly-dead state. Symptoms range from a ~30s window of dead DNS to a permanent outage that only full `recover.sh` or `toggle.sh` fixes. The root cause is that nullexit has **\*\*no daemon watching macOS sleep/wake or network-state-change events\*\*** the existing DNS Watcher inside `toggle.sh` only runs in-process and is suspended along with the rest of the user shell on lid close, so the gateway cannot self-heal after either event.

1046

1047 **\*\*Diagnosis:\*\*** Two distinct failure modes share the same root gap.

1048

1049 **\* \*\* (A) Lid close sleep wake.**  `caffeinate -i` (used by `toggle.sh`) blocks **\*idle\*** sleep but **\*\*does not** block forced sleep from lid close**\*\***. Closing the lid hard-suspends Colima VM + every container + every Tailscale-quit-shells-except-tailscaled. On wake the VM clock needs ~5-30s to resync via NTP, during which TLS cert validation fails: Cloudflare WARP `162.159.192.1:2408` rejects the handshake gluetun's healthcheck (interval: 2s, retries: 15) eventually flips unhealthy Docker `unless-stopped` restarts the `warp` container (~30s gap). `mDNSResponder`'s in-memory cache still holds pre-sleep entries (stale A records for tailnet hosts) so the first dozen DNS queries after wake time out. `tailscaled` on the host often keeps its stale DERP relay preference and routes packets through a dead relay for another 30-60s.

1050 **\* \*\* (B) Wi-Fi roam / loss in elevators, captive portals.**  WARP is a UDP tunnel to Cloudflare's anycast edge. Carrier NATs and captive-portal networks silently invalidate the UDP binding on roam and on hotspot-bound re-association. macOS `networksetup` **\*should\*** preserve DNS settings on the same Wi-Fi service across a roam, but most captive-portal networks DHCP-replace DNS to the captive-portal resolver **\*\*during the captive-portal dance itself\*\***, before the user has authenticated so even DNS hijack survives a clean roam and quietly breaks on captive-portal handoff.

1051

1052 **\*\*Resolution.\*\*** A single **\*\*launchd LaunchAgent\*\*** (`com.nullexit.wake-recovery`) runs a long-running shell daemon (`scripts/watcher.sh`) that listens for both event sources and, on each fire, executes `bash recover.sh --post-wake` a **\*lighter-touch\*** recovery mode that's the opposite of the existing `--nuclear` semantics: it keeps the gateway live while still refreshing every stale subsystem.

1053

1054 **\*\*Deep dive: Apple Power Management Primer (for the unfamiliar).\*\***

1055 macOS exposes several relevant surfaces; the right primitive depends on what you actually need to detect. None of these are obvious and none of them have a standard splash screen in the docs.

1056

1057 **\* \*\* `caffeinate <flags>`**  creates I/O Kit power assertions via `IOPMAssertionCreateWithName`. It does NOT block forced sleep (lid close, low-battery shutdown, sleep timer firing on a per-set Energy Settings policy). Flags:

1058 **\* `i`** PreventIdleSleep (the only one `toggle.sh` uses; protects against the OS going idle while the user is active but a clamshell close still hard-puts it to sleep)

1059 **\* `s`** PreventSystemSleep (only honored on AC power; the user may be on battery)

1060 **\* `u`** DeclareUserActive (resets the idle timer every 30s by default; equivalent to keeping the cursor moving)

1061 **\* `d`** PreventDisplaySleep

1062 **\* `m`** PreventDiskIdleSleep

1063 **\* There is NO `l` (lid-block) flag in any modern macOS.** Lid close is a kernel-firmware-level event that ignores user-space assertions.

1064 **\* To prove any of this on live hardware:** `pmset -g assertions` lists every active assertion with type / process / reason / timeout.

1065 **\* \*\* `pmset`**  is the read-side of the same system:

1066 **\* `pmset -g log | grep -E 'Sleep|Wake|DarkWake`** shows the recent timeline.

1067 **\* `pmset -g history`** is the per-day summary.

1068 **\* `pmset -g assertions`** is the live assertion list (matches `i s d u` to processes).

1069 **\* \*\* `launchd`**  does NOT have a native "on-wake" hook. **\*\* Not in any version of macOS as of Sonoma.** The canonical recipe is `StartCalendarInterval` with a sub-minute cadence and let launchd queue missed intervals for replay-on-wake, or use a third-party daemon (Sleepwatcher). For our use case a long-running shell daemon listening to the unified log is more responsive than a polling agent.

1070 **\* \*\* `com.apple.system.powermanagement.\*` Darwin notifications**  (`willSleep`, `didWake`, `didDim`, `didUndim`) are the C-level signal. From a shell, they are best reached through the unified log with `log stream --predicate` see the watcher implementation.

```

1071 * scutil n.watch is the canonical CLI for live-following network-state changes. Pairs with n.add State:/
    Network/Global/IPv4 to get a single tap that fires on any IPv4 link/route/DNS/address change across all
    interfaces. This is what scripts/watcher.sh's network-change listener uses.
1072 * com.apple.networkChange Darwin notification is the C-level equivalent but has no shell-friendly
    listener; use scutil instead.
1073 * SCNetworkReachability is the Objective-C API used by SOCKS5/HTTP clients to detect route changes per
    target. Never a fit for shell scripts.
1074
1075 **The Fix (component-by-component).**
1076
1077 1. recover.sh --post-wake (new flag, non-destructive by design). Adding --post-wake to the existing
    nuclear recovery script inverts the semantics. The default mode still tears down Tailscale, resets DNS to
    empty, stops caffeinate, runs docker compose down, power-cycles Wi-Fi. The new mode does:
1078 * Skip Tailscale disconnect (we WANT it to keep the mesh connection)
1079 * Re-hijack DNS to the gateway Tailscale IP read from .gateway_ip (instead of resetting to empty)
    applies to both ACTIVE_SERVICE and ENO_SERVICE because the system resolver is scoped to en0
1080 * Refresh the exit-node preference with tailscale up --reset --ssh=true --accept-dns=false --exit-node
    ="$TS_IP" --exit-node-allow-lan-access=true (the exact flags toggle.sh START uses). This re-asserts
    DERP mapping without dropping the mesh
1081 * Inspect docker inspect --format '{{.State.Health.Status}}' nullexit-warp-1 and only docker compose
    up -d --force-recreate warp if gluetun is unhealthy this single targeted recreate nudges the UDP NAT
    binding back to Cloudflare without disturbing Tailscale or AdGuard
1082 * Run the sharingd-reset step (existing fix from toggle.sh START path; see 15.11) so AirDrop doesn't freeze
    on the new IP lease
1083 * Verify the gateway with dig +tcp @127.0.0.1 -p 5354 google.com (AdGuard via TCP through Colima's SSH
    tunnel) and a 4-second tailscale ping to the gateway Tailscale IP, instead of the default mode's direct-
    to-1.1.1.1 check
1084 * Crucially: skip sudo -v because the LaunchAgent has no TTY to prompt on; all privileged calls use sudo -n
    and lean on /etc/sudoers.d/nullexit NOPASSWD entries (networksetup, dnsmasq, scacheutil,
killall, pkill).
1085 2. toggle.sh writes and clears /tmp/nullexit-gateway-active.marker (an ISO-8601 UTC timestamp) at the
    bottom of the START path and the top of the STOP path / cleanup_handler. Two new helpers: write_gateway_active_marker
    and clear_gateway_active_marker. The watcher uses this file as its only
    signal that the gateway is live: if toggle.sh was stopped (or never started), the watcher no-ops. This
    prevents waking-up-only-on-counterfactual events from churning DNS.
1086 3. scripts/watcher.sh (long-running daemon, sibling-style source file). Two backgrounded pipelines:
1087 * Wake listener: log stream --predicate 'composedMessage CONTAINS[c] "didWake" OR composedMessage
    CONTAINS[c] "Wake from Sleep" OR composedMessage CONTAINS[c] "Waking from" OR composedMessage CONTAINS[c] "DarkWake"
    --style compact --info piped through while read; do run_recover "WAKE: ..."; done <[c]
    makes the predicate case-insensitive (the unified log writes phrases in mixed case across versions).
1088 * Network listener: (echo "n.add State:/Network/Global/IPv4"; echo "n.watch"; while true; do sleep
    86400; done) | scutil 2>/dev/null piped through a case match on n.state / SCEventUpdate.
1089 * run_recover checks the marker, debounces via /tmp/nullexit-watcher.last-recovery (default 10s,
    configurable via NULLEXIT_DEBOUNCE_SECONDS), and shells out to bash recover.sh --post-wake. The 10s
    debounce prevents a single wake event (which typically emits 2-3 log lines) from triggering multiple
    recoveries.
1090 * trap ... TERM INT kill 0 cleanly tears down the process group on launchctl bootout so the sleep
    86400 inside the scutil pipe also dies.
1091 * wait blocks indefinitely while both pipelines feed it.
1092 4. launchd/com.nullexit.wake-recovery.plist (LaunchAgent in the user domain runs as the console user at
    every login without needing sudo).
1093 * Label: com.nullexit.wake-recovery
1094 * ProgramArguments: ["/bin/bash", "<absolute-path-to-your-nullexit-install>/scripts/watcher.sh"]
1095 * RunAtLoad: true starts immediately on launchctl load (no need to log out and back in)
1096 * KeepAlive: { SuccessfulExit: false, Crashed: false } don't restart on clean SIGTERM exit (so launchctl bootout
    doesn't get stuck in a relaunch loop). Also don't restart on hard crash: we observed
    that macOS Sonoma's sleep/wake suspend-resume path accumulates orphan watcher.sh PIDs because SIGCONT
    does not reliably clean up the prior instance's listener grandchildren, and a KeepAlive.Crashed=true-
    driven relaunch actively races with the still-live prior process. The script enforces single-instance on
    its own via a flock-style PID-file lock, so a relaunch-on-crash would be skipped by the lock and produce
    0 listener coverage until the live instance happened to die. Trade-off: a real crash leaves the user
    without auto-recovery until they run launchctl load -w manually acceptable because (a) gateway still
    works without auto-recovery, (b) a relaunch wouldn't fix whatever caused the crash, and (c) the gateway-
    recovery itself is observably broken (no DNS, no exit-node) if it does go down.
1097 * ProcessType: Background hint to App Nap not to suspend us.
1098 * StandardOutPath/StandardErrorPath: /tmp/nullexit-watcher.{out,err}.log separates launchd-level
    diagnostics from the script's own /tmp/nullexit-watcher.log (which exec >>> redirects).
1099
1100 **Install (one-time).** The plist lives in the repo at launchd/com.nullexit.wake-recovery.plist for
    version control. The installed (live) copy must land in ~/Library/LaunchAgents/ so launchd picks it up on
    the user session.
1101
1102 ```bash
1103 # Copy the plist to the user-launchd location (run from repo root)
1104 cp ./launchd/com.nullexit.wake-recovery.plist \
1105 ~/Library/LaunchAgents/
1106
1107 # Substitute __NULLEXIT_HOME__ with your local install absolute path.
1108 # The plist uses WorkingDirectory + a relative program arg, so this
1109 # single substitution is the only place it references your machine.
1110 sed -i '' "s|__NULLEXIT_HOME__|$(pwd)|" \

```

```

1111 ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1112
1113 # Validate the XML
1114 plutil -lint ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1115
1116 # Load + persist this user session + every future login
1117 launchctl load -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1118
1119 # Confirm it actually started
1120 launchctl list | grep nullexit
1121 ```
1122
1123 To **stop** the watcher (without uninstalling):
1124
1125 ```bash
1126 launchctl bootout gui/$UID/com.nullexit.wake-recovery
1127 # or:
1128 launchctl unload -w ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1129 ```
1130
1131 To **uninstall** completely:
1132
1133 ```bash
1134 launchctl bootout gui/$UID/com.nullexit.wake-recovery 2>/dev/null || true
1135 rm ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
1136 ```
1137
1138 **Why this fix was preferred over alternatives.**
1139
1140 * **Polling with `StartCalendarInterval`** would have given ~30-60s detection latency per wake. The unified-log
1141   stream is sub-second.
1142 * **Sleepwatcher** would have covered wake events but not network changes. We'd still need a separate `scutil`
1143   listener, and introducing a third-party dependency for what amounts to ~150 lines of shell is unjustified.
1144 * **Forcing `caffeinate -l`** doesn't exist and the closest equivalent (`caffeinate -s`) only works on AC power
1145   . The whole point of the gateway is to stay usable on battery while traveling lid close must NOT silently
1146   kill the gateway.
1147 * **A kernel extension** is impossible (no entitlement) and out of scope.
1148
1149 **Reproduction recipe.** Test both events in isolation to confirm the fix actually closes the gap.
1150
1151 * **(A) Lid-only repro**: with the gateway up, run in a terminal: `pmset displaysleepnow && sleep 30 && pmset
1152   wake` (without physically closing the lid). Watch the wake wake-fires the watcher post-wake routine runs
1153   login should still work in <5s after wake. Compare to baseline pre-fix: pre-fix the gateway would be dead
1154   for 30-90s.
1155 * **(B) Roam-only repro**: with the gateway up: `networksetup -setairportpower off && sleep 30 && networksetup
1156   -setairportpower on` (or simply walk out of Wi-Fi range and back). The captive-portal WILL repave DHCP DNS
1157   for a few seconds; the watcher fires on the second `State:/Network/Global/IPv4` change (after the Wi-Fi re-
1158   associates) and the post-wake routine re-hijacks DNS within 2-3s.
1159
1160 **Operative files.**
1161
1162 * `recover.sh` added arg parsing; wrapped destructive sections behind `POST_WAKE`; added re-hijack DNS /
1163   refresh exit-node / force-recreate-if-unhealthy path.
1164 * `toggle.sh` added `write_gateway_active_marker` / `clear_gateway_active_marker` called at the END of START
1165   and TOP of STOP / cleanup_handler.
1166 * `scripts/watcher.sh` new long-running daemon.
1167 * `launchd/com.nullexit.wake-recovery.plist` launchd LaunchAgent descriptor.
1168 * `~/Library/LaunchAgents/com.nullexit.wake-recovery.plist` the installed copy (one-time copy).
1169
1170 **Honest assessment (read before testing).** This fix has two asymmetric halves:
1171
1172 * **Network / roam / captive-portal path RELIABLE.** `scutil n.watch` on `State:/Network/Global/IPv4{4,6}`
1173   catches every Wi-Fi roam, captive-portal DHCP-rebind, hotspot handoff, and VPN change with zero missed
1174   events. To validate, drop Wi-Fi for 30 s and re-add it; `/tmp/nullexit-watcher.log` will record a `NET:`
1175   trigger and `recover.sh --post-wake` will run within ~10 s of the network coming back.
1176 * **Lid close / wake path BEST-EFFORT on macOS Sonoma.** Apple *suspends* LaunchAgents process-state during
1177   forced sleep and *resumes* them on wake; it does NOT re-launch them every wake. As a result: (a) `log
1178   stream` is a live subscription events emitted during the agent's suppressed window are DROPPED, not
1179   buffered, so the wake event that prompted the resume is invisible to `log stream` on resume; (b) the "fire
1180   on startup" guard runs once on `launchctl load -w` and after a `KeepAlive` crash-recovery, NOT on every
1181   successive wake; (c) `subsystem == "com.apple.powermanagement"` will catch FUTURE emissions (display dim/
1182   restore, manual sleep/wake) but not the lid-close/open wake directly.
1183
1184 In practice the lid-wake gap is short, because `toggle.sh`'s existing DNS Watcher already polls every 5 s and
1185 re-hijacks DNS (so DNS recovers within 5 s of any roam or wake), `tailscaled` self-heals via DERP fallback
1186 within 3060 s, and `gluetun` reconnects via its healthcheck retry within 30 s. The user-visible pain on lid
1187 open is ~30 s of wonky DNS, not a permanent outage.
1188
1189 **Test the ROAM path FIRST.** If ROAM works in your testing, the lid-wake gap is acceptable. If lid-wake
1190 handling is critical, the two clean follow-ups are:

```

```

1166 1. **Sleepwatcher** (one canonical install): `brew install sleepwatcher`; drop `recover.sh --post-wake` into
    `~/sleep` and `~/wake` so Sleepwatcher invokes it directly on every wake without the launchd propagation
1167 2. **Move the watcher to a LaunchDaemon** (root) and use `IOPMSchedulePowerEvent` via a tiny C wrapper not
    portable to shell.
1168
1169 A truly reliable shell-only solution to "wake events from inside a LaunchAgent" does not exist on macOS Sonoma
    +.
1170
1171 **Empirical lessons from deployment.** Two findings came up while deploying this fix end-to-end; recording so
    the next operator doesn't lose an evening to each.
1172
1173 1. **Bash function-call-before-definition gotcha.** While introducing the gateway-active marker helpers, the
    defensive `clear_gateway_active_marker` call was inserted at the very top of `toggle.sh` (right after `set
    -e`), but the function definition lived near `PID_FILE="/tmp/nullexit-cafeinate.pid"` ~90 lines down. Bash
    parses top-to-bottom and resolves names at first invocation, so the call site exited with code 127 (`
    clear_gateway_active_marker: command not found`). The gateway stayed unreachable from the START path even
    though `bash -n` passed (the parser doesn't catch order-of-resolution errors). Rule: define functions
    BEFORE any block that calls them. Verify with `grep -n 'name()'s*{'` followed by `grep -n 'name\s*${'`
    the latter must appear on a line number AFTER the former.
1174 2. **networksetup -setairportpower off+on` is not a faithful roam reproduction.** Synthetic Wi-Fi radio-
    cycling (`sudo networksetup -setairportpower Wi-Fi off` for 30 s, then back on) is expected to produce a `
    NET:` trigger in `scutil n.watch` but produced ZERO during testing because `setairportpower` powers the
    radio at the driver level while leaving the network SERVICE in the `up` state from scutil's perspective,
    so no `n.state State:/Network/Global/IPv4` event is generated. A real roam (macOS briefly loses the AP and
    re-associates into a different one) DOES fire the event because the link actually changes.
1175 Better reproduction recipes in priority order:
1176 * Walk between two real SSIDs via `networksetup -switchtolocation` (or actual physical station movement)
    guaranteed to fire the scutil event.
1177 * Force a DHCP rebind on the same SSID with `sudo ipconfig set en0 DHCP` triggers `State:/Network/Global/
    IPv4` to update.
1178 * Trust the existing 30-second DNS Watcher inside `toggle.sh` for the DNS path: it re-hijacks DNS
    automatically. The post-wake watcher adds clear value mostly for tailscale exit-node re-assertion and warp
    gluetun force-recreate, both of which self-heal within ~30 s anyway.
1179
1180 ##### 15.5.2 July 8, 2026: Network Watcher Event Mismatch, Sudo-in-Background Crash, and Loop-Prevention
1181 **Symptom:**
1182 1. Switching Wi-Fi or waking the Mac from sleep did not trigger a post-wake recovery (`recover.sh --post-wake`)
    automatically.
1183 2. The network connection would eventually get completely sinkholed/blocked. If the user tried to restart or
    wait, it did not self-heal.
1184 3. During the recovery attempt, checking the public IP on the host would return the unencrypted campus/ISP IP (
    starting with `132.`) instead of the encrypted tunnel IP.
1185
1186 **Root Cause:**
1187 - **Bug 1 (Network Watcher Mismatch):** In `scripts/watcher.sh`, the network change listener watched `State:/
    Network/Global/IPv4` changes via `scutil n.watch` but matched them against `*n.state*|*SCEventUpdate*`. On
    macOS, the actual output is `changedKey [0] = State:/Network/Global/IPv4`. Because the pattern did not
    match, all network switches were silently ignored.
1188 - **Bug 2 (Sudo Caching Background Crash):** When the background WARP Watcher eventually detected the broken
    tunnel (after 6 failures/30s), it triggered the default nuclear `recover.sh` (with `POST_WAKE=false`).
    Because there was no active TTY in the background context, the `sudo -v` caching check aborted instantly
    with a terminal required error. Due to `set -e`, `recover.sh` died immediately before resetting DNS,
    disabling the packet-filter killswitch, or stopping containers, leaving the host network completely
    sinkholed.
1189 - **Bug 3 (Post-Wake Infinite Loop):** Once the watcher patterns were fixed, `recover.sh --post-wake` triggered
    successfully on network changes. However, because the script deleted the default routing entries at the
    start, spent ~15 seconds rebuilding 88 DERP relay bypass routes, and re-added the default routes at the end
    , this execution time exceeded the watcher's 10-second debounce threshold. Furthermore, the watcher
    recorded the `start` timestamp in the debounce file. Consequently, the route changes made by the script
    itself triggered the watcher again immediately after completion, trapping the system in an infinite loop of
    recovery runs. During this loop, the VPN routes were constantly deleted, resulting in the host leaking
    traffic through the real ISP interface (an IP starting with `132.`) for 15s out of every cycle.
1190
1191 **Resolution:**
1192 - **Fixed Watcher Matching:** Updated `scripts/watcher.sh` to match `*changedKey*` and `*State:/Network/Global/
    IPv*` to correctly catch network changes.
1193 - **Bypassed Background Sudo:** Added `--auto`/`--non-interactive` flags to `recover.sh` and added a TTY check
    `[ -t 0 ]` to skip interactive `sudo -v` authentication when running headlessly in the background. Updated
    the WARP Watcher call in `toggle.sh` to pass `--auto`.
1194 - **Prevented Infinite Loops:** Modified `scripts/watcher.sh` to write the **completion** timestamp of `recover
    .sh` to `DEBOUNCE_FILE` (instead of only the start timestamp). This ensures all buffered network config
    events generated during route reconstruction are evaluated against the end-of-run timestamp and
    successfully debounced, breaking the infinite loop.
1195 - **Centralized & Timestamped Logs:** Redirected `watcher.sh` stdout/stderr from `~/tmp/nullexit-watcher.log` to
    the repository's main `output.log` and updated `scripts/common.sh`'s `step`, `ok`, `warn`, `fail`, and `
    die` helpers to prefix logs with a local `[YYYY-MM-DD HH:MM:SS]` timestamp.
1196
1197 ##### 15.5.3 Incident Post-Mortem: Startup Pre-Flight Check Tailscale Race (July 10, 2026)
1198 **Symptom:** Immediately after startup via `toggle.sh`, the pre-flight connectivity check `[1/3] Gateway
    reachable via Tailscale` returned `FAIL`, forcing `toggle.sh` to skip the full exit-node activation and

```



```

1199     fall back to SOCKS5/local DNS proxy mode.
1200 **Root Cause:** A timing race condition during startup. In `toggle.sh`, the host joins the mesh using `
    tailscale up --reset` (Phase A) right before performing the pre-flight check. Because `tailscale up` resets
    and reconfigures the host's `tailscaled` daemon, the local daemon must fetch the latest netmap
    asynchronously from the coordination servers. This network map propagation and route establishment takes a
    few seconds. Running `tailscale ping` immediately after `tailscale up` returns (with no delay) causes the
    check to fail because the local daemon has not yet discovered the gateway node `100.100.21.8`.
1201
1202 **Fix (July 10, 2026):** Upgraded pre-flight Check 1 in both `toggle.sh` and `scripts/toggle-linux.sh` to use a
    5-attempt retry loop with a 1-second delay between checks. This allows up to 5 seconds for local tailnet
    routing paths to settle, ensuring the gateway is correctly reached and a pong is received before deciding
    whether to enable exit-node routing.
1203
1204 #### 15.5.4 Incident Post-Mortem: Pre-Flight Check False Negatives due to VPN Firewall STUN Blocks (July 10,
    2026)
1205 **Symptom:** Immediately after a cold boot or full restart of the gateway, the pre-flight connectivity check
    `[1/3] Gateway reachable via Tailscale` failed, causing `toggle.sh` to skip the full exit-node
    configuration and fall back to SOCKS5/local DNS proxy mode. However, manually running `tailscale ping
    100.100.21.8` immediately after startup succeeded in 1ms.
1206
1207 **Root Cause:** An asynchronous Tailscale handshake delay caused by firewall restrictions. Inside the container
    network namespace, the `warp` container (gluetun) implements a strict firewall that blocks all outbound
    UDP traffic to the public internet except to the designated WireGuard endpoint. Because of this, the `
    tailscale` container's STUN requests to public STUN servers are blocked, causing the local daemon to report
    `UDP is blocked, trying HTTPS`. Without UDP STUN capability, Tailscale is forced to fall back to a slower
    DERP-assisted handshake (relayed via HTTPS/TCP). This negotiation and direct-path punch-through takes
    roughly 30 to 35 seconds to establish on cold boots (after a full `tailscale up --reset`). Since the pre-
    flight check only waited 15 seconds (15 attempts with 1-second sleeps), the check timed out and aborted
    before the path completed.
1208
1209 **Fix (July 10, 2026):** Increased the pre-flight ping retry limit from 15 to 45 attempts (45 seconds) in both
    `toggle.sh` and `scripts/toggle-linux.sh`. This provides sufficient time for the DERP-assisted handshake to
    complete on cold boots, while still passing immediately (in 1-2 seconds) on warm starts or once the path
    is established.
1210
1211 #### 15.5.5 Incident Post-Mortem: Docker Image Build DNS Failures on Immediate Restart (July 10, 2026)
1212 **Symptom:** When executing `toggle.sh --restart`, the script successfully stops the gateway but then
    immediately fails during startup during `docker compose build` with DNS resolution errors:
1213 `failed to do request: Head "https://registry-1.docker.io/...": dial tcp: lookup registry-1.docker.io on
    192.168.5.1:53: no such host`.
1214
1215 **Root Cause:** A race condition between physical link negotiation and virtual machine startup. The STOP path
    of the restart command invokes `cleanup_network_state` which power-cycles the host's physical Wi-Fi
    interface (`en0`) to flush stale routing states. Immediately after, `toggle.sh` starts the gateway boot
    sequence. Because `toggle.sh` did not verify the status of the physical interface, it proceeded to boot
    Colima and build Docker containers while `en0` was still negotiating a DHCP lease. Consequently, the host (
    and by extension, the VM bridge) had no active internet gateway/DNS path, causing Docker's registry check
    to fail.
1216
1217 **Fix (July 10, 2026):**
1218 1. Created a unified `wait_for_dhcp_settle` helper function in `scripts/common.sh` that polls `ipconfig
    getpacket` and `ifconfig` until a valid IP and router are assigned. To handle typical macOS Wi-Fi
    reassociation and DHCP negotiation latency (which commonly takes 15-20 seconds after a power-cycle), this
    loop was configured to poll up to 60 times (30 seconds total).
1219 2. Injected a call to `wait_for_dhcp_settle` at the beginning of the `START` action in `toggle.sh` (right
    before checking Colima status) to block container builds until the host's physical network is online.
1220 3. Refactored `recover.sh` to reuse the unified `wait_for_dhcp_settle` function.
1221
1222 #### 15.5.6 Watcher.sh Suppressed Exit Code Bug (Fixed)
1223 **Observation:** The `scripts/watcher.sh` script is a long-running daemon that invokes `recover.sh --post-wake`
    when it detects system wake or network roam events. Previously, it contained the following bug:
1224 ```bash
1225 bash "$RECOVER" --post-wake
1226 local exit_code=$?
1227 ...
1228 return $exit_code || true
1229 ```
1230 The `|| true` mistakenly swallowed the actual `$exit_code` returned by `recover.sh`, causing `run_recover()` to
    always return `0` (success), masking any underlying failures in the wake recovery process.
1231
1232 **Fix:** The `|| true` statement was removed from the return line:
1233 ```bash
1234 return $exit_code
1235 ```
1236 Since `watcher.sh` explicitly omits `set -e` (to prevent pipe breakages from killing the daemon), removing `||
    true` safely allows the underlying exit code to propagate without risking daemon termination. The watcher
    daemon was then reloaded (`launchctl kickstart -k`) to pick up the script changes in memory.
1237
1238 #### 15.5.7 Network Readiness Race Condition on `--restart` (Fixed)

```

```

1239 **Problem:** When running `./toggle.sh --restart`, the script tears down the gateway and immediately begins the
      startup sequence. On systems with Wi-Fi (or any wireless interface), the OS takes a moment to re-associate
      with the network after the teardown's DNS/routing changes. If the startup sequence ran before the OS had a
      default route, `get_active_service` would return nothing, routing bypass targets would be wrong, and DNS/
      WARP setup would silently fail resulting in a partially configured gateway that self-heals only once the `
      watcher.sh` re-runs.

1240
1241 **Root cause:** The original code had no gate before the first network-dependent call (`ACTIVE_SERVICE=$(
      get_active_service)` at the top of the start branch). This was a pure timing race.

1242
1243 **First fix (Wi-Fi specific):** A polling loop on `ipconfig getifaddr en0` was added to wait for an IPv4
      address on `en0` before proceeding. This worked but was brittle it only handled Wi-Fi and would fail for
      Ethernet, USB-C adapters, or iPhone USB tethering.

1244
1245 **Generalized fix:** Replaced the `en0` check with `route get default | grep 'gateway:'`. This detects a valid
      default gateway on *any* interface, covering all connection types. The loop polls every 2 seconds with a
      60-second hard timeout (then proceeds with a warning rather than hanging forever). The output prints the
      detected gateway IP so failures are easy to diagnose in logs.

1246
1247 ```bash
1248 # In toggle.sh, before ACTIVE_SERVICE=$(get_active_service)
1249 while true; do
1250     if route get default 2>/dev/null | grep -q 'gateway:.'; then
1251         _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
1252         echo " Network ready (default gateway: $_gw)."
1253         break
1254     fi
1255     ...
1256 done
1257 ```

1258
1259 **Why not ping?** Pinging an IP (e.g. `1.1.1.1`) during restart is unreliable because DNS may still be hijacked
      from the previous session, and the exit node routing may not yet be cleared, causing the ping to loop back
      through the tunnel. `route get default` is a pure local kernel query no packets sent.

1260
1261 ### 15.6 WARP Tunnel Liveness & Host Leaks

1262
1263 #### 15.6.1 Host-Side Traffic Leaking Past the Exit Node (with Remote Clients Routing Correctly)

1264 **Symptom.** `toggle.sh` reports success. Remote tailnet clients (e.g. an S24 phone) correctly egress through
      Cloudflare WARP and see a Cloudflare IP on `whatismyip.com`. But on the **HOST Mac running the gateway,
      `whatismyip.com` reports the underlying physical ISP's ASN (e.g. `campus_isp` for a university campus Wi-Fi)
      . The gateway container itself is healthy; the leak is specifically on the host.

1265
1266 **Root causes.** Three distinct mechanisms can each produce this exact symptom on the host alone, while leaving
      remote clients unaffected. They are mutually rankable so the diagnostic picks the active one
      deterministically.

1267
1268 * **(A) SOCKS5 fallback lane active.** `toggle.sh`'s three Phase-B pre-flight checks (`tailscale ping` AdGuard
      via `dig +tcp` WARP internet, toggle.sh ~lines 750-820) sometimes flake during the first minute. When ANY
      of the three fails, the script sets `SKIP_EXIT_NODE=true` (toggle.sh:849) and instead activates the SOCKS5
      fallback path (toggle.sh:894). The SOCKS5 fallback hijacks DNS to `127.0.0.1` (so AdGuard filtering still
      works) but **modern browsers on macOS do NOT honor the system SOCKS5 proxy** they ignore `networksetup -
      setsocksfirewallproxy`. Result: DNS gets ad-filtered but actual HTTP/S traffic exits direct through the
      campus ISP. This branch is invisible from a remote tailnet client because the client's tunnel is unaffected
      .

1269
1270 * **(B) IPv6 leak over campus Wi-Fi.** The Tailscale exit node and WARP tunnel are both IPv4-only (gluetun +
      post-rules.txt` drop all IPv6 forwarded traffic; see 15.12 IPv6 Exit-Node Leak). But many university and
      enterprise APs are dual-stack the host happily accepts AAAA DNS responses and sends IPv6 traffic straight
      out `en0`. macOS happily encrypts that traffic through Tailscale ONLY if Tailscale has IPv6 advertised;
      with an IPv4-only exit node, IPv6 traffic skips Tailscale entirely. This branch is invisible from a phone
      because iOS/Android Tailscale apps actively sinkhole IPv6 in their VPN profile, but the macOS host's
      `tailscaled` does not.

1271
1272 * **(C) Route-table freeze and `--accept-routes` reset after `tailscaled` wake/toggle.** Running `tailscale up
      --reset` clears the exit-node preference and reverts `--accept-routes` back to its default (`false`).
      Without explicitly passing `--accept-routes=true`, `tailscaled` will not attempt to install the default
      route through the `utun*` tunnel interface even if the exit node preference is reconciled. Additionally,
      macOS occasionally freezes and fails to apply route-table changes (see 15.11 Standalone Tailscale Daemon
      Freeze). In both cases, `netstat -rn` shows the default route still pointing to the physical Wi-Fi gateway
      (e.g. `172.17.0.1`) instead of the Tailscale `utun*` interface, causing host traffic to exit direct while
      remote clients route correctly.

1273
1274 **Why remote clients don't see this.** Each remote phone/laptop uses its OWN Tailscale tunnel to reach the
      container's `100.x.x.x:443` over the mesh they're end-to-end encrypted regardless of how the host Mac
      itself is configured. Only the HOST's own egress sockets (HTTP requests from the host's browser, curl, etc
      .) are affected.

1275
1276 **Resolution.** A single script: `scripts/diagnose-host-leak.sh`. It runs the 8 checks, classifies into one of
      A / B / C / OK, prints a verdict + a ready-to-run fix command on stdout, AND writes the full report to
      `host-leak-diagnostic-<UTC>.txt` so the user can paste the file back instead of scrolling-and-copying from

```



```

the terminal. With `--fix`, the matched remediation is applied and the two checks that could change (warp +
ip6) are re-verified. With `--watch` (or `--watch 30`), it runs the full baseline diagnostic then enters
a continuous loop checking warp/IPv6/default-route every N seconds, alerting only on state changes logs to
`host-leak-watch.log`.
1277
1278 ```bash
1279 # Diagnose only:
1280 bash scripts/diagnose-host-leak.sh
1281
1282 # Diagnose + apply matched fix + re-verify:
1283 bash scripts/diagnose-host-leak.sh --fix
1284 ```
1285
1286 **What it prints.** A 7-section report (`tailscale status`, SOCKS5 state, `cdn-cgi/trace` warp=, IPv6 leak
probe, host default route, last toggle.sh output.log hints, resolved gateway Tailscale IP), followed by a
classification block (Scenario: A | B | C | OK) with the exact command(s) the user needs to paste into
their terminal they don't have to figure out which `tailscale up` flags match their environment.
1287
1288 **What it does NOT do.** It does not touch `toggle.sh` or `toggle.sh`'s pre-flight logic, does not modify `.env`
(except `--fix` mode appends `GATEWAY_BYPASS_PING=true` when fixing Scenario A), and does not restart
Docker. It is a read-only diagnostic with an opt-in remediation flag. Safe to run on a healthy gateway the
worst case is a 30-line report that says "OK".
1289
1290 **Why cdn-cgi/trace over whatismyip.com for the verdict.** `whatismyip.com`'s ISP database routinely mislabels
campus IPv6 ranges as residential ISPs (that's why the local ISP's ASN appears even when you're routing
through Cloudflare). `cdn-cgi/trace` is hosted by Cloudflare ITSELF and reports `warp=on|off` plus the
exact edge colo serving the request. It is the only public endpoint that can truthfully confirm whether the
WARP tunnel is in use.
1291
1292 **Common false alarms.**
1293
1294 * **`warp=on` but whatismyip.com still shows the local ISP.** `whatismyip.com`'s ISP-detection database is
unreliable on academic networks. Trust `cdn-cgi/trace`'s `warp=on` over `whatismyip.com`'s ISP string. Use
`https://api.ipify.org` (returns just the IP, no ISP DB lookup) as a third confirmatory check.
1295 * **`tailscale status` shows the host online but exit node preference is missing.** `tailscale up --reset --
exit-node=` (empty value) clears any stale preference; `tailscale up --reset --exit-node=<100.x.x.x>` re-
establishes it. The diagnostic prints the exact command with the resolved TS_IP filled in.
1296 * **Both IPv6 leak AND route-freeze flags set simultaneously.** Scenario B outranks Scenario C fixing IPv6
makes IPv4 the only viable egress, after which the route-freeze usually self-resolves within ~30s (or one
more `toggle.sh` cycle).
1297
1298 **Deep dive: Host Leak Probe Continuous Sub-Second Egress Monitor (`scripts/host-leak-probe.sh`).**
1299
1300 > ** STATUS: REMOVED (historical).** `scripts/host-leak-probe.sh` and the `HOST_LEAK_PROBE` env flag no longer
exist the file is deleted/untracked and nothing launches it (verified 2026-07-15). Its fail-closed role
was superseded by the **PF kill-switch** (`enable_killswitch` in `common.sh`) plus the in-container **WARP
Watcher** (15.6.2); the only surviving host-leak tool is the on-demand `scripts/diagnose-host-leak.sh`. The
text below is kept for design rationale only treat every present-tense claim (auto-launch, PID file, `
HOST_LEAK_PROBE=false` toggle) as past tense.
1301
1302 **What it is.* `scripts/host-leak-probe.sh` is a lightweight background daemon (~1.4 MB RSS, ~01% CPU) that
polls `https://www.cloudflare.com/cdn-cgi/trace` every 300ms directly from the host via `curl` not via `
docker compose exec`. This is a fundamentally different vantage point from the WARP Watcher (15.6.2): the
WARP Watcher asks the WARP *container* whether it sees `warp=on`; the Host Leak Probe asks what a real
browser request leaving the host's physical NIC looks like. The two blind spots it covers that the WARP
Watcher cannot:
1303
1304 1. **Host-side flash-leaks** a Cloudflare anycast edge re-route, a browser reusing a stale pooled connection
across a tunnel state change, or a split-second routing gap during a Gluetun healthcheck restart (see
15.12.13 Routing Fix 30-second polling acknowledged 14s window).
1305 2. **Host egress while container is healthy** the container can report `warp=on` while the host's own traffic
exits direct (the three-scenario leak class documented above in 15.6.1).
1306
1307 **Lifecycle.* Started by `toggle.sh`'s START path (`start_host_leak_probe`) immediately after `
start_warp_watcher`. Controlled by:
1308
1309 | Signal path | What happens |
1310 |---|---|
1311 | `toggle.sh` STOP | `stop_host_leak_probe()` SIGTERM + PID file cleanup |
1312 | `toggle.sh` `cleanup_handler` (error/INT/TERM/HUP) | Same |
1313 | `recover.sh` (default, non `--post-wake`) | 6d block kill by PID file |
1314 | `recover.sh --post-wake` | Left running the gateway is still live |
1315
1316 PID file: `/tmp/nullexit-host-leak-probe.pid`. Toggle with `HOST_LEAK_PROBE=false` in `.env` to disable at next
gateway start.
1317
1318 **Output format.* All events are written to `output.log` (repo root) with UTC timestamps. Only state *changes*
are logged the probe is silent during normal healthy operation.
1319
1320 ```
1321 [09:10:26] LEAK warp=off ip=45.23.1.1 prev_warp=on prev_ip=104.28.246.50

```

```

1322 [09:10:26] ROTATE warp=on ip=104.28.248.11 prev_ip=104.28.246.50
1323 [09:10:26] HOST-PROBE failed/timeout (count=1)
1324 ```
1325
1326 curl transport errors (`TLS handshake failed`, `Connection refused`, etc.) are also appended to `output.log`
1327 via `2>>output.log` nothing is silenced to `/dev/null`.
1328
1329 *Grep cheatsheet.*
1330
1331 ```bash
1332 grep LEAK output.log          # real warp=off events from the host
1333 grep ROTATE output.log        # Cloudflare anycast edge rotations (harmless)
1334 grep 'HOST-PROBE' output.log   # curl failures / timeouts
1335 ```
1336
1337 *Resource budget.* ~1.4 MB RSS, ~0% CPU at rest, ~3 HTTPS requests/second to Cloudflare. Safe to leave running
1338 for hours. Back off the polling interval (`bash scripts/host-leak-probe.sh 1.0`) if you observe probe-
1339 failure pile-ups indicating Cloudflare rate-limiting.
1340
1341 *Detection vs prevention.* This script detects leaks; it does not prevent them. A sub-300ms flash-leak can
1342 still slip between two polls undetected. If `grep LEAK output.log` shows confirmed leak events, the next
1343 step is a kill-switch (prevention) an `iptables` rule on the host's `en0` that drops non-Tailscale, non-
1344 local egress whenever the gateway is active. That is a separate engineering task not yet implemented.
1345
1346 ##### 15.6.2 In-Flight WARP Tunnel Liveness Monitor & Statistical Auto-Shutdown
1347
1348 **Why.** nullexit's core promise is that your IP always appears as Cloudflare. If the WARP tunnel silently dies
1349 due to a UDP NAT timeout, a Cloudflare edge rotation, a Colima VM clock skew causing TLS cert rejection,
1350 or any of a dozen other failure modes the host silently falls back to egressing through the physical ISP.
1351 The user sees the gateway as "up" (containers running, Tailscale connected, DNS hijacked) but their real IP
1352 is exposed. This is the exact class of failure that Ingo Blechschmidt documented in his chilling 2024
1353 Linux kernel post-mortem: for over two years, LUKS disk encryption on suspend was silently broken because a
1354 refactored kernel syscall returned success while doing nothing ([source](https://mathstodon.xyz/@iblech
1355 /116769502749142438)). The lesson: **a security mechanism that fails silently is worse than no mechanism at
1356 all** it creates a false sense of safety that prevents the user from taking corrective action.
1357
1358 **Design principle.** The WARP Watcher (`start_warp_watcher()` in `toggle.sh`) is a background daemon that
1359 polls `cdn-cgi/trace` every 30 seconds and applies a **statistically calibrated threshold** before
1360 triggering any safety response. This threshold distinguishes transient network blips (which self-heal
1361 within seconds) from genuine outages (which require intervention). A single `warp=off` reading is
1362 meaningless Docker might be slow, the network might be congested, a container healthcheck might be
1363 restarting. But 6 consecutive off readings (30 continuous seconds of downtime) has vanishingly low
1364 probability of being a false positive: even at an unrealistically high 10% false-positive rate per poll, P
1365 (6 consecutive) = 0.1 0.0001%. In practice, the per-poll false-positive rate is far lower than 10%, making
1366 the real probability astronomically small.
1367
1368 **What it does.**
1369
1370 1. **Always logs.** Every state transition (onoff, offon) is timestamped to `output.log`. The user can `grep
1371 WARP output.log` to audit the tunnel's entire lifetime. This is the "never fail silently" mandate.
1372
1373 2. **Tracks consecutive failures.** A counter increments on each `off` reading and resets to 0 on any `on`
1374 reading. This ensures the watcher never fires on intermittent flapping.
1375
1376 3. **Auto-shuts down at threshold.** When the counter reaches `WARP_FAIL_THRESHOLD` (default 6, configurable in
1377 `.env`), the watcher declares a statistically significant outage and runs `recover.sh` the nuclear
1378 recovery script that disconnects Tailscale, stops Docker containers, resets DNS to `1.1.1.1`, flushes
1379 routes, and power-cycles Wi-Fi. This instantly reverts the host to normal (un-encrypted) internet,
1380 guaranteeing no traffic leaks through a dead tunnel while the user thinks they're protected. A trigger file
1381 at `/tmp/nullexit-warp-shutdown-triggered` prevents double-firing.
1382
1383 4. **Exits cleanly.** After shutdown, the watcher exits (the gateway is dead, there's nothing left to monitor).
1384 `stop_warp_watcher()` is also called by `toggle.sh`'s STOP path and `cleanup_handler` to terminate the
1385 daemon cleanly during normal teardown.
1386
1387 **Configuration.** Set `WARP_FAIL_THRESHOLD=3` in `.env` for a more aggressive 15s timeout, or `
1388 WARP_FAIL_THRESHOLD=12` for a conservative 60s window. The variable is parsed from `.env` using the same `
1389 grep` pattern as `GATEWAY_MSS` and `GATEWAY_BYPASS_PING`.
1390
1391 **Log sample.** After a real outage (or a forced test via `docker compose pause warp`):
1392
1393 ```
1394 [2026-07-03T21:15:17Z] WARP DOWN cdn-cgi/trace reports warp=off
1395 [2026-07-03T21:15:47Z] WARP SHUTDOWN 6 consecutive failures (threshold=6). Running recover.sh to kill the
1396 gateway and restore normal internet. Your IP is NO LONGER Cloudflare.
1397 [2026-07-03T21:15:52Z] WARP SHUTDOWN recover.sh completed, watcher exiting. Your IP is now your real ISP IP.
1398 ```
1399
1400 Or, if WARP self-recovers before the threshold:
1401
1402 ```
1403 [2026-07-03T21:15:17Z] WARP DOWN cdn-cgi/trace reports warp=off

```

```

1369 [2026-07-03T21:15:32Z] WARP RECOVERED warp=on after 3 consecutive off readings (threshold was 6)
1370 ```
1371
1372 **Relationship to other monitors.** The WARP Watcher is the **innermost safety layer** it runs as a child of `
toggle.sh` and only while the gateway is active. It complements (does not replace) the `scripts/watcher.sh`
daemon (15.5.1, which handles sleep/wake and Wi-Fi roam) and `scripts/diagnose-host-leak.sh` (15.6.1,
which is a manual diagnostic tool). The WARP Watcher is the only component that actively verifies end-to-
end tunnel liveness and takes autonomous action when it fails.
1373
1374 #### 15.6.3 Incident: The Overnight Silent IP Leak (July 2026)
1375 **The Problem:** The host leaked traffic over its raw, unencrypted `eth0` IP continuously for roughly 8 hours,
completely bypassing the VPN. The `toggle.sh` state and the container liveness checks all reported the
gateway as healthy and active.
1376
1377 **The Investigation:**
1378 1. **Sleep/Wake Checks:** We initially suspected macOS sleep cycles broke the routing table. A check of `pmset
-g log` confirmed the laptop literally never slept (0 sleep events recorded since boot), ruling out all
sleep/wake code paths.
1379 2. **Keepalive Expiry:** We discovered `docker-compose.yml` was missing `
WIREGUARD_PERSISTENT_KEEPA_LIVE_INTERVAL`. Without a keepalive, standard router NAT tables (which typically
garbage collect UDP after 30 to 120 seconds of silence) quietly expired the WireGuard port mapping
overnight.
1380 3. **Container State Masking:** Because WireGuard is stateless, the `warp` container didn't immediately crash.
It stayed "Up", which meant Docker's healthchecks and our `toggle.sh` liveness checks never noticed the
tunnel inside it was totally dead.
1381
1382 **The Solution:**
1383 1. Added `WIREGUARD_PERSISTENT_KEEPA_LIVE_INTERVAL=25s` (the WireGuard standard to beat standard 30s NAT
timeouts) to keep the UDP tunnel permanently pinned open through the router.
1384 2. Injected a **Fail-Closed Kill-Switch** into `routing-fix.sh`: A 30-second loop now actively verifies tunnel
health by running `curl` over `tun0` to Cloudflare's trace endpoint. If it fails 3 consecutive times, it
immediately rewrites the default route to `blackhole`, severing all traffic instead of letting it leak
through `eth0`.
1385 3. Added a listener in `watcher.sh` that detects this fail-closed event and instantly pushes a native macOS UI
notification to the desktop, alerting the user to manually intervene.
1386
1387 > [!NOTE]
1388 > **Why dual failure systems?**
1389 > The system now relies on two independent failure handlers that cover each other's blind spots:
1390 > 1. **Automated Self-Healing (WARP Watcher + `recover.sh`):** Triggered when the `warp` container logs `warp=
off`. This happens when the server actively kicks the client. Because the container *knows* it is
disconnected, it triggers `recover.sh` to safely reboot Docker and reconnect.
1391 > 2. **Fail-Closed Kill-Switch (`routing-fix.sh` + `watcher.sh`):** Triggered when a silent network failure
occurs (e.g., NAT timeouts). The container incorrectly believes it is connected (`warp=on`), so auto-
recovery never fires. Instead of auto-recovering (which risks falling back to raw Wi-Fi mid-process while
the user isn't looking), this kill-switch actively blackholes the internet and forces a manual intervention
via a desktop notification.
1392
1393 ### 15.7 Tailscale P2P, SNAT & Control Plane
1394
1395 #### 15.7.1 macOS SNAT Endpoint Poisoning & The P2P Blackhole (July 10, 2026)
1396 **Symptom:** When a Tailscale client (like an S24) connected to the exact same Wi-Fi network as the macOS
gateway, it completely lost internet connectivity. However, if the S24 connected to a Windows PC's hotspot
on the same network, the tunnel worked flawlessly.
1397
1398 **Root Cause:** Claude's critique correctly dismantled the port collision theory: the gateway daemon wasn't
failing to bind, and macOS wasn't load-balancing ports. The real culprit was **macOS SNAT Source Port
Mangling poisoning Tailscale's path discovery**.
1399
1400 Here is the exact step-by-step of the "infinite loop" blackhole:
1401 1. **The P2P Handshake:** The S24 and the Mac are on the same Wi-Fi (`192.168.1.x`). The S24 discovers the Mac'
s local IP and sends a UDP hole-punch packet to `192.168.1.5:41641`. Colima successfully forwards this to
the gateway container.
1402 2. **The Bypass Rule:** The gateway replies. Because `routing-fix.sh` forces all Tailscale UDP traffic out `
eth0` to bypass WARP, the reply goes to the macOS host.
1403 3. **macOS SNAT Mangling:** macOS routes the reply out `en0` to the S24. Because the packet crosses from the
Colima VM bridge to the Wi-Fi interface, macOS applies Source NAT (SNAT). Critically, macOS **randomizes
the source port** (e.g., changes it from `41641` to `58291`).
1404 4. **Endpoint Poisoning:** The S24 receives the authenticated Tailscale packet from `192.168.1.5:58291`.
Tailscale's `magicsock` is designed to handle NAT dynamically, so it assumes the Mac has roamed to port
`58291`! The S24 updates its endpoint and starts sending all its encrypted internet traffic to
`192.168.1.5:58291`.
1405 5. **The Blackhole:** macOS has no port forward for `58291`. It drops 100% of the S24's outbound data packets.
The connection is completely dead.
1406
1407 **Why did the Hotspot work?*** When the S24 connected to the PC hotspot, it was behind the PC's NAT. The S24
sent the packet, PC NATted it, and the Mac replied (mangled to `58291`). When the reply hit the PC, the PC'
s strict NAT dropped it because the source port (`58291`) didn't match the destination port it originally
sent to (`41641`). Because the mangled packet was dropped, the S24 never poisoned its endpoint! It safely
gave up on P2P, fell back to the 100% reliable TCP DERP relay, and the internet worked perfectly.
1408

```

```

1409 **Fix & Resolution (July 10, 2026):**
1410
1411 **Phase 1 Port disambiguation:** Changed the gateway container's Tailscale listen port from the default
    `41641` to `41642` across `docker-compose.yml`, `post-rules.txt`, and `routing-fix.sh`. This prevents any
    bind conflict with the macOS host's native Tailscale daemon.
1412
1413 **Phase 2 RFC1918 DROP rules:** Updated `routing-fix.sh`'s `add_tailscale_p2p_bypass()` to explicitly drop all
    outgoing Tailscale UDP packets destined for private subnets (`192.168.0.0/16`, `10.0.0.0/8`,
    `172.16.0.0/12`) when `TAILSCALE_ALLOW_LAN_P2P=false`. This surgically kills the local P2P hole-punch
    attempt *before* it can trigger the macOS gvproxy SNAT mangling. The S24 receives a clean failure and
    immediately locks onto the public DERP relay with 0% packet loss.
1414
1415 **Phase 3 Automatic network detection:** Because `TAILSCALE_ALLOW_LAN_P2P=false` breaks P2P on trusted
    hotspots that genuinely don't have AP isolation, a `.env` toggle and a full auto-detection system were
    implemented:
1416
1417 - **`.env` toggle:** `TAILSCALE_ALLOW_LAN_P2P=true/false` static fallback, used only if auto-detection cannot
    run.
1418 - **`watcher.sh` auto-detection (`detect_lan_p2p_mode`):** On every network change, `watcher.sh` runs two
    checks:
1419 1. **WPA2-Enterprise detection:** `airport -I` is used to read the current Wi-Fi `link auth` field. If the
    auth type does not contain `psk` (e.g., it is plain `wpa2` = 802.1x), the network is classified as
    enterprise and P2P is forced off. *(Note: the `airport` binary was removed in macOS Sonoma 14.4 the
    detection now uses `system_profiler SPAirPortDataType`; see 15.11 for the full migration.)*
1420 2. **AP Isolation probe:** For WPA2-Personal networks, a short `ping` is sent to the default gateway. If the
    gateway responds (it always should on a real home/hotspot network), P2P is allowed. If not, AP isolation is
    inferred and P2P is forced off.
1421 - The result (`true` or `false`) is written to `.lan_p2p_detected` in the repo root.
1422 - **`routing-fix.sh` dynamic reading:** `add_tailscale_p2p_bypass()` re-reads `.lan_p2p_detected` on every 30-
    second loop iteration and atomically adds or removes the RFC1918 DROP rules. **No container restart is
    needed when switching networks.**
1423
1424 **Known Unknowns (Pending Verification):**
1425 - The gvproxy SNAT source port mangling theory is mechanistically sound and backed by observed behavior (
    gvproxy's UDP proxy changelog has documented edge cases in this exact code path), but has not yet been
    confirmed via `tcpdump -i en0 udp portrange 40000-60000` while triggering a P2P handshake from the phone. A
    packet capture cross-referenced with `tailscale ping` endpoint reports on the phone side would confirm the
    theory definitively.
1426 - Claude's recommendation: run `tcpdump -i en0 udp port 41642 or portrange 40000-60000` on the Mac while
    triggering the handshake from the phone to see whether the reply leaves with a different source port than
    `41642`.
1427 - **Client-Side VPN Cycling (Roaming Recovery):** An observation (July 10, 2026) suggests that when a hung DNS
    probe state is triggered by roaming between Access Points (APs) on a WPA2-Enterprise network, simply
    toggling the VPN connection (Tailscale) off and on locally on the client phone immediately resolves the
    issue. You do not need to cycle the phone's physical Wi-Fi. This reinforces the theory that roaming on
    Enterprise networks leaves a stale/poisoned P2P endpoint in the phone's Tailscale magicsock state, which
    gets instantly flushed upon a local VPN restart. (Status: Possible but not formally verified).
1428
1429 #### 15.7.2 July 6, 2026: Packet Filter (pfctl) Defaulting Local LAN Rules to TCP-Only
1430 **Symptom:** Devices on the exact same local Wi-Fi network as the gateway were experiencing ~500ms latency to
    the gateway instead of the expected ~80ms.
1431
1432 **Root Cause:** When writing the `pf.conf` rules to whitelist local LAN traffic (`pass out quick on $ext_if to
    { 192.168.0.0/16, ... }`), the rule omitted an explicit protocol. The macOS `pfctl` compiler implicitly
    applies state tracking to all `pass` rules. However, because state tracking (`keep state`) natively
    evaluates TCP sessions, `pfctl` automatically appended the `flags S/SA` modifier to the compiled rule in
    the kernel. This silently restricted the whitelist to **TCP only**, causing all local UDP traffic to be
    dropped by the default-deny kill-switch. Since Tailscale relies on UDP port 41641 for direct P2P
    connections, the local devices were unable to hole-punch and were forced to fall back to routing traffic
    through a remote Tailscale DERP relay, massively inflating latency.
1433
1434 **Resolution:** Split the single implicit rule into explicit `proto tcp`, `proto udp`, and `proto icmp` rules
    in `scripts/pf.conf` to guarantee the macOS compiler allows all three core transport protocols on the local
    subnet without mutating the rule into a TCP-only lock.
1435
1436 #### 15.7.3 Incident Post-Mortem: PF Kill-Switch Bricks Host Internet in SOCKS5 Fallback Mode (July 10, 2026)
1437 **Symptom:** When the pre-flight checks fail (e.g. because host Tailscale was unauthenticated/logged out), the
    script falls back to SOCKS5 proxy mode but the host immediately loses all internet connectivity and
    tailscaled reports `You are logged out... connect: no route to host`. Even running `sudo tailscale up` to
    log in fails because the connection to the control plane times out.
1438
1439 **Root Cause:** An architectural mismatch. The macOS Packet Filter (PF) Kill-Switch works at the IP level: it
    blocks all outbound traffic on the physical interface (`en0`) except to explicitly whitelisted VPN
    endpoints, expecting all other traffic to go through the virtual VPN interface (`utun*`). In SOCKS5
    fallback mode, the exit-node (`utun*`) is NOT enabled; instead, only specific applications (like browsers)
    use the localhost SOCKS5 proxy, while all other host traffic (including the `tailscaled` daemon's UDP
    packets) continues to go through `en0`. Because the script still enabled the PF Kill-Switch during SOCKS5
    fallback, it blocked all non-SOCKS5 traffic on `en0`, completely bricking the host's internet and locking
    the Tailscale daemon out of the control plane.

```

1441 **\*\*Fix (July 10, 2026):\*\*** Modified `toggle.sh` to remove the `enable_killswitch` call from the SOCKS5 fallback path. The firewall kill-switch is now strictly reserved for exit-node VPN mode, ensuring SOCKS5 fallback leaves the host's direct internet route open for system daemons to authenticate.

1442

1443 **#### 15.7.4 Incident Post-Mortem: Host Tailscale Logouts due to Aggressive reset Flags (July 10, 2026)**

1444 **\*\*Symptom:\*\*** Every time the user runs `toggle.sh` or `recover.sh` which disconnects/reconnects the host client, the host client randomly gets logged out, forcing them to re-run `sudo tailscale up` to authenticate.

1445

1446 **\*\*Root Cause:\*\*** An overly aggressive configuration reset. The host-side `tailscale up` calls (used in `disconnect_tailscale_host` and the startup/enable exit node paths) all passed the `--reset` flag. On macOS, `--reset` resets the entire configuration state and authentication token cache of the daemon. Since the script does not pass a `TS_AUTHKEY` to the host's commands (which are meant for the user's private interactive client), the `--reset` flag effectively logged the user out of their tailnet, requiring interactive re-login.

1447

1448 **\*\*Fix (July 10, 2026):\*\*** Removed the `--reset` flag from all host-side `tailscale up` invocations in `toggle.sh` and `scripts/common.sh`. This ensures that exit-node state transitions (enabling/disabling) are handled safely while preserving the host client's authenticated session.

1449

1450 **#### 15.7.5 Incident Post-Mortem: macOS Tailscale Flag Requirement Abort (July 10, 2026)**

1451 **\*\*Symptom:\*\*** When the `toggle.sh` stop trap or the roaming watcher triggered, the gateway failed to shut down or disconnect properly. The `output.log` showed the error: `Error: changing settings via 'tailscale up' requires mentioning all non-default flags. To proceed, either re-run your command with --reset...`

1452

1453 **\*\*Root Cause:\*\*** In a previous attempt to stop host-side logouts (15.7.4), we removed the `--reset` flag from all `tailscale up` invocations. However, if the Tailscale daemon on macOS is already running with non-default flags (like `--ssh` or `--accept-routes`), invoking `tailscale up` with only a subset of flags (like `--exit-node=`) triggers a strict safety check in the Tailscale CLI that aborts the command completely. Because the command aborted, exit-node routing remained stuck.

1454

1455 **\*\*Fix (July 10, 2026):\*\*** Migrated all specific preference changes in `toggle.sh` and `scripts/common.sh` from `tailscale up` to a two-step sequence:

1456 1. `tailscale up` (with no flags) to safely ensure the interface is brought up using the current authenticated state.

1457 2. `tailscale set --exit-node=...` to modify the specific target preference cleanly without triggering the missing flags error or requiring `--reset`.

1458

1459 **### 15.8 Docker / Colima Lifecycle**

1460

1461 **#### 15.8.1 Colima VM Memory / Swap Thrashing Latency**

1462 **\*\*Symptom:\*\*** After running nullexit flawlessly for over 24 hours straight to test stability, the 0.5GB VM RAM configuration began experiencing severe network latency on the Tailscale mesh later in the day.

1463

1464 **\*\*Root Cause:\*\*** Services (like AdGuard, Tailscale, Docker logs) slowly accumulate cache and state over extended periods. Under the tight 512MB physical RAM limit, this slow creep forced the Linux kernel to aggressively swap memory to the 512MB SSD swap file. The constant disk I/O thrashing blocked process execution, causing DNS resolutions and packet forwarding to delay significantly (making the internet feel "very slow").

1465

1466 **\*\*Proof (Empirical Observation):\*\*** While in this OOM-thrashing state, direct DNS queries through the Tailscale mesh (`dig @100.100.21.8 google.com`) would intermittently hang or take several seconds to resolve. After restarting with the 600MB configuration, the exact same query immediately dropped to a lightning-fast **\*\*41 ms\*\*** response time. (Note: Querying the mapped host port via `127.0.0.1:5354` will always time out due to Gluetun's strict leak-prevention firewall blocking non-VPN ingress traffic).

1467

1468 **\*\*Fix:\*\*** Increased VM base physical RAM to 600MB (`--memory 0.6`) and reduced the swap file to 400MB. This provides enough native RAM headroom for long-running state caches without forcing heavy I/O swapping. Subdomain deduplication still reduces blocklists by ~60%. Memory profiles (`light`/`medium`/`heavy`) let users trade off coverage for memory.

1469

1470 **#### 15.8.2 Missing iptables in routing-fix Container**

1471 **\*\*Root Cause:\*\*** `alpine:3.20` does not have iptables installed. Commands failed silently with `2>/dev/null || true`.

1472

1473 **\*\*Fix:\*\*** Updated `docker-compose.yml` to `apk add --no-cache iptables iproute2` before `scripts/routing-fix.sh`.

1474

1475 **#### 15.8.3 DOCKER\_GW Variable Parsing Bug**

1476 **\*\*Root Cause:\*\*** `ip route show default` printed two lines; `awk '{print $3}'` picked up both, causing route commands to fail.

1477

1478 **\*\*Fix:\*\*** Added `head -1` to the parsing chain.

1479

1480 **#### 15.8.4 Hard Reboots Leave Stale Docker Socket in Colima VM**

1481 **\*\*Problem:\*\*** If the macOS host machine is subjected to a hard reboot, sudden power loss, or a kernel panic, the Colima VM running in the background is abruptly killed. Upon next boot, running `toggle.sh` resulted in the contradictory output:

1482 `Colima is already running.` followed by `Cannot connect to the Docker daemon at unix:///var/run/colima/default/docker.sock.`

1483

1484 **\*\*Root Cause:\*\*** The hard crash prevents the inner Docker daemon from executing its graceful shutdown routines. It leaves behind a corrupted `docker.sock` file inside the VM. On boot, the outer VM spins up (hence "



already running"), but the inner Docker daemon sees the stale socket, assumes another instance is running, and crashes.

**\*\*Fix:\*\*** Added an Auto-Healing routine directly into `toggle.sh` (Step 4b). The script now explicitly tests the Docker daemon with `docker ps`. If the daemon is unresponsive despite Colima reporting as active, it assumes a stale socket and automatically runs `colima restart` in the background to cleanly rebuild the VM and socket before proceeding.

**#### 15.8.5 Cloudflare WARP 24-Second Startup Delay (UDP Session Rate-Limiting)**

**\*\*Context:\*\*** When running `toggle.sh` to turn the gateway ON, `docker compose up -d` often hangs for exactly 24-25 seconds waiting for the `warp` container to become `healthy`. Users might mistakenly blame the Go `rule-compiler` processing 400k+ rules for this delay, as Docker prints `rule-compiler Exited 24.3s` at the end of the wait.

**\*\*Root Cause:\*\*** The `rule-compiler` actually finishes its work in <2 seconds. The 24-second blockage is entirely due to Cloudflare WARP's edge server rate-limiting. WireGuard is a stateless protocol with no "disconnect" packet. When the toggle is turned OFF, the old UDP session state remains active on Cloudflare's backend for 20-30 seconds. When the toggle is instantly turned back ON, Docker starts a new WireGuard connection from a new ephemeral UDP source port but uses the same static cryptographic key (from `.env`). Cloudflare detects this as a potential replay attack or key-sharing abuse, and intentionally drops all packets (rate-limits) for the new connection until the old session's ghost state fully expires (~20-25 seconds).

**\*\*Result:\*\*** Gluetun's healthcheck fails repeatedly every 2 seconds until Cloudflare lifts the penalty box, at which point traffic flows and Docker unblocks the rest of the startup sequence. This is a known, expected behavior of reusing static keys rapidly on Cloudflare's network and cannot be bypassed.

**#### 15.8.6 July 4, 2026: Multi-stage Docker Builds for Go Ports & Teardown Hang Fix**

**\*\*Symptom:\*\*** The Python implementations of `logger.py` and `sync-rules.py` were slow and required a heavy python runtime inside Alpine containers. Also, `routing-fix` container teardown was hanging for 30s during `toggle.sh` shutdown.

**\*\*Root Cause:\*\***

- Python is slower than Go and installing it dynamically via `apk add` added overhead and complexity to the containers.
- Docker containers ignore `SIGTERM` if the init process doesn't handle it, causing a full 30s timeout on shutdown.

**\*\*Resolution:\*\***

- Ported `logger` and `sync-rules` to Go using Goroutines for massive concurrent speedups.
- Implemented **\*\*Multi-stage Docker builds\*\*** (using `golang:1.22-alpine` as builder) to compile the static binaries inside Docker. The final containers are raw Alpine with zero dependencies.
- Added `--build` to `docker compose up -d` in `toggle.sh` to ensure updates are consistently applied on boot.
- Added `stop\_grace\_period: 1s` to `routing-fix` in `docker-compose.yml` which eliminates the 30-second teardown hang instantly.
- Implemented a cryptographic integrity checker (`scripts/crypto.sh`) using HMAC-SHA256 and `NULLEXIT\_SEED` in `.env` to prevent tampering of bash scripts. (See also 15.9.4.)

**#### 15.8.7 Colima `start` Post-Provision Hang (~2 min) & the Docker-Readiness Poll (July 14, 2026)**

**\*\*Symptom:\*\*** `toggle.sh` took ~197s of wall-clock time to bring the gateway up, and almost the entire delay lived in a single step: `colima start`. Timestamps in `output.log` showed the VM reaching `READY` and creating the `colima` Docker context in ~8-12 seconds, after which `colima start` simply \*did not return\* for ~110 more seconds until the `run\_with\_timeout 120` watchdog `SIGKILL`ed it. The rest of the boot (containers, routing) then proceeded normally.

**\*\*Root Cause:\*\*** A colima-internal hang that occurs \*after\* `Successfully created context colima` is printed. It is **\*\*mode-independent\*\***. It was initially misattributed to `--network-address` / `--network-mode bridged` needing the `socket\_vmmnet` helper, but a controlled test with plain user-mode networking (`colima start --memory 0.6 --vm-type vz`, no `--network-address`) reproduced the identical ~120s hang (04:40:34 04:42:34, watchdog-killed) while Docker was fully usable the entire time. Crucially, killing the hung `colima start` process does **\*\*not\*\*** stop the VM and does **\*\*not\*\*** corrupt state `colima status` correctly reports `running` afterwards and `docker` keeps working via the `colima` context. The ~110s was therefore pure dead time waiting on a CLI that had already done its real work.

**\*\*Fix:\*\*** `scripts/common.sh` now provides `colima\_start\_until\_ready()`. It launches `colima start "\$@"` in the background and polls `docker --context colima info` every 3s. The moment Docker answers (~12s in practice) it stops waiting, reaps the (usually still-hung) starter with `kill`, runs `docker context use colima` to guarantee the default context is set, and returns 0. If Docker never answers within a 90s cap it returns non-zero and `toggle.sh` continues to the existing Step 4b `docker ps` auto-heal (15.8.4). Both Colima start call sites in `toggle.sh` the LAN-P2P bridged/shared path and the user-mode fast path now go through this helper instead of `run\_with\_timeout 120 colima start`.

**\*\*Result:\*\*** The Colima phase dropped from ~120s to ~12s and total `toggle.sh --restart` wall time from ~197s to ~106s, with the double-tunnel and direct hostcontainer P2P confirmed intact after the change. The key insight: treat "Docker is reachable" not "the `colima start` CLI returned" as the definition of \*started\*.

**#### 15.8.8 Toggle-On Startup Optimizations: Rule-Compiler Off the Critical Path (July 14, 2026)**

**\*\*Context:\*\*** A genuine cold toggle-on (VM bounced for RAM, WARP down) measured **\*\*95s\*\***. Profiled per-phase with the now-fixed `DEBUG\_TRACE` (15.11.8-A): Colima cold start **\*\*19s\*\*** (near floor), `docker compose up` **\*\*45s\*\***, tailscale connect+poll **\*\*21s\*\***, final checks **\*\*7s\*\***.

**\*\*Two wrong hypotheses first (documented so they aren't re-tried):\*\*** (1) that BuildKit `--build` was the ~30s cost it wasn't; every layer was already `CACHED`, so gating `--build` saved only ~2s. (2) that `docker compose exec` was slow in the poll loop it isn't (0.18s steady-state vs 0.06s for `docker exec`); the ~5s/



iteration was `tailscale status` **\*\*blocking during cold connect\*\*** and hitting the `run\_with\_timeout 5` cap.

**\*\*Real bottleneck (evidence-backed):\*\*** the 45s `compose up` was **\*\*not\*\*** WARP. On a genuine cold start WARP goes healthy fast (tor starts right with it no 15.8.5 same-key penalty; that only bites rapid offon). The long pole is the **\*\*rule-compiler re-deduping ~340k rules inside the 600MB VM (~28s), every toggle\*\*** `adguardhome` started 28s after warp because it waited on `rule-compiler: service\_completed\_successfully`, even though all 16 remote lists loaded from cache unchanged.

**\*\*Fix hash-gated, off the critical path (kept in-container):\*\***

- `rule-compiler` is now a **\*\*compile--profiled\*\*** service, excluded from `docker compose up`. `adguardhome` and `routing-fix` no longer `depends\_on` it they boot instantly on the `compiled\_rules.txt` / `ip\_blocklist.ipset` already in `./adguard/work`.
- `toggle.sh` runs it on-demand via `docker compose run --build --rm rule-compiler`, **\*\*gated on `compute\_rules\_hash()`\*\*** = sha256 of `black\_list.txt` + `white\_list.txt` (the lists the user controls). Recompile only when that hash changes, when an artifact is missing, or when the compiled file is older than `RULES\_MAX\_AGE\_DAYS` (default 7, so the remote blocklists still refresh occasionally). A compile failure is **\*\*non-fatal\*\*** we keep the previous compiled rules rather than bricking the tunnel (ad-block the kill-switch).
- Also trimmed the tailscale-connect poll timeouts (status `5s3s`, `ip -4` `10s5s`) to reduce overshoot.

**\*\*Why NOT run the compiler natively on the Mac host (the tempting idea):\*\*** rejected on purpose. There is no Go on the host, so "native" means either `brew install go` (a host toolchain dependency that drifts the project from \*portable-secure-on-any-device\* toward \*faster-on-mine\*) or committing a prebuilt arch-specific **\*\*binary blob\*\*** you'd then have to trust/sign. Worse, it would make the **\*\*host\*\*** write into the `./adguard/work` bind-mount violating the deliberate **\*\*single allowed writer is the rule-compiler container\*\*** invariant (hostcontainer UID/perms mismatches corrupt AdGuard's BoltDB store). Staying in-container preserves both properties; the gate delivers the speed.

**\*\*Result:\*\*** unchanged toggle `compose up` **\*\*45s 9s\*\***, total **\*\*--restart` \*\*\*95106s 63s\*\***, with AdGuard still serving the full 342,519 rules (it reuses the existing compiled file) and the double-tunnel + direct hostcontainer P2P intact. Editing `black\_list.txt`/`white\_list.txt` flips the hash one-off recompile new rules take effect. Both lists were added to the `crypto.sh` signed set (they're security inputs), so editing them requires a re-sign which is the same moment you'd want the recompile. Local state files `build\_hash` / `rules\_hash` are gitignored.

### ### 15.9 Permissions, Crypto & Security Hardening

#### #### 15.9.1 Sudo Credential Caching Removed

**\*\*Problem:\*\*** The script used `sudo -v` at startup to cache sudo credentials, plus a background `SUDO\_KEEPER\_PID` loop refreshing them every 60 seconds. This required interactive password entry even with NOPASSWD sudoers configured, because `sudo -v` itself prompts.

**\*\*Fix:\*\*** Removed the entire `sudo -v` + `SUDO\_KEEPER\_PID` section. All privileged commands now use `sudo -n` (non-interactive), which works silently because the user has NOPASSWD rules in `/etc/sudoers.d/nullexit` for the specific commands needed (`networksetup`, `python3`, `dscacheutil`, `killall`, `pkill`, `route`, `ifconfig`).

#### #### 15.9.2 Background Sudo Failures

**\*\*Context:\*\*** The `--post-wake` script is intentionally designed to skip the `sudo -v` interactive prompt because it is launched by `watcher.sh` running via `launchd` in the background (which has no TTY and cannot accept password input). Thus, it relies entirely on `sudo -n` (non-interactive sudo) for all its internal routing changes.

**\*\*Problem:\*\*** If the user has not configured the `/etc/sudoers.d/nullexit` file properly, any automated background wake event will silently fail: `sudo -n` commands drop out, the WARP bypass routes fail to apply, and the `warp` container loses internet connectivity.

#### #### 15.9.3 Tailscale Hijacking the Default Route (PHYSICAL\_GW Bug)

**\*\*Problem:\*\*** When `--post-wake` runs, it clears the routing table using `sudo route -n flush` to ensure a clean slate after the Wi-Fi card roams or wakes up. Because it skips bouncing the Wi-Fi interface (to make the reconnect faster), macOS's DHCP client does not automatically rebuild the physical default route. To counteract this, the script must manually restore the physical default route. Originally, the script tried to capture the physical gateway via `route get default`. However, because the gateway is already running and Tailscale is active, the default route is often already pointing to `utunX`. When this happens, `route get default` does not output a `gateway:` field (it only outputs the `interface:`), causing the captured `PHYSICAL\_GW` variable to be empty.

When `PHYSICAL\_GW` is empty, the physical default route is never restored, causing the `en0` interface to be isolated from the physical network. The WARP bypass routes then fall back to targeting `--interface en0` (instead of routing via the gateway), which breaks WireGuard's remote connections completely.

**\*\*Fix:\*\*** The script now bypasses the OS routing table entirely and directly queries the hardware DHCP lease (`ipconfig getpacket en0`) to reliably extract the physical router IP, ensuring the physical route is correctly restored even when Tailscale is active.

#### #### 15.9.4 July 4, 2026: Consolidation of Core Bash Scripts (Refactoring)

**\*\*Symptom:\*\*** Over 200 lines of bash logic were directly duplicated across `toggle.sh` and `recover.sh` (e.g. `restart\_tailscaled\_daemon`, `stop\_sleep\_prevention`, WARP bypass routes). This caused drift when one script was updated without the other.

**\*\*Root Cause:\*\*** Historical separation of responsibilities allowed `recover.sh` to grow complex alongside `toggle.sh`.

**\*\*Resolution:\*\***

- Massively refactored both scripts by extracting all duplicated networking logic, PID lifecycle management, and environmental configuration down to `scripts/common.sh`.

```

1551 - Specifically, the `stop_sleep_prevention`, `stop_dns_watcher`, `stop_host_leak_probe`, and `stop_warp_watcher`
1552       functions were collapsed into a single `stop_pidfile_daemon()` abstraction.
1553 - Proxy disable loops were abstracted to `disable_all_proxies()`.
1554 - The `162.159.192.1/193.1` WARP edge IPs are now dynamically resolved from `.env` instead of hardcoded.
1555
1556 ##### 15.9.5 Threat Model: Local Proxy Discovery
1557 The Tor SOCKS proxy and Honey-Port Tripwire only bind to `127.0.0.1` (loopback). The real security boundary
1558 here is "can an unprivileged local process reach the loopback interface," not "does the malware know the
1559       port number."
1560
1561 The port-randomization and the Honey-Port trap only catch blind or generic port-scanners that do not already
1562 know where to look. They **do not protect** against a process that directly reads the `/tmp/nullexit-ports.
1563 env` file, greps the `toggle.sh` memory space, or runs `ls -i` to inspect open socket bindings.
1564
1565 Because modifying file ownership on `/tmp/nullexit-ports.env` via `chown`/`chmod` to restrict read-access
1566 introduces an unavoidable Time-Of-Check to Time-Of-Use (TOCTOU) race condition (the file is created user-
1567 owned before privilege escalation, meaning file descriptors can be snatched), it is deliberately left as a
1568 standard user-owned file.
1569
1570 The actual, proper mitigation for preventing a malicious local process from reaching the proxy is not hiding
1571 the portit is running a host-level outbound firewall with strict localhost/loopback filtering enabled (e.g
1572 ., LuLu 4.3.0+ or Little Snitch). These firewalls gate access by cryptographic process identity (binary
1573 signature) at the kernel/network-extension level, rather than by port secrecy.
1574
1575 **Verifying Localhost Filtering (The Rogue Binary Test).** To empirically validate that the host firewall
1576 properly intercepts local proxy hijacking, you can compile and execute a completely unknown, un-whitelisted
1577 "rogue" binary to attempt a connection to the proxy port:
1578
1579 ```c
1580 // lulu-test.c
1581 #include <stdio.h>
1582 #include <stdlib.h>
1583 #include <arpa/inet.h>
1584
1585 int main(int argc, char *argv[]) {
1586     int port = atoi(argv[1]);
1587     int sock = socket(AF_INET, SOCK_STREAM, 0);
1588     struct sockaddr_in server = { .sin_family = AF_INET, .sin_port = htons(port) };
1589     server.sin_addr.s_addr = inet_addr("127.0.0.1");
1590
1591     printf("Rogue binary attempting connection to 127.0.0.1:%d...\n", port);
1592     if (connect(sock, (struct sockaddr *)&server, sizeof(server)) < 0) {
1593         printf("Connection BLOCKED by firewall.\n");
1594         return 1;
1595     }
1596     printf("Connection SUCCESSFUL (Firewall bypassed!)\n");
1597     return 0;
1598 }
1599 ```
1600
1601 Compile with `gcc -o lulu-test lulu-test.c`. When executed against the `$TOR SOCKS_PORT`, a properly configured
1602 host firewall (with "Allow Loopback" disabled in LuLu Preferences) will immediately pause the socket
1603 execution and generate a desktop alert for the unrecognized binary signature. This physically stops
1604 targeted local malware from exfiltrating data through the Tor tunnel, proving the threat model mitigation.
1605
1606 ##### 15.9.6 Unlocking Mode-Locked Output Files Without `chmod` (`unlock-files.sh`)
1607
1608 **The one-command fix:** `bash scripts/unlock-files.sh` replaces every known locked inode in the project with
1609 a fresh writable one. No `chmod`. Covers `.env` (mode `000`), `adguard/work/userfilters/compiled_rules.txt`
1610 , `ip_blocklist.ipset`, `cache/*.txt`, `cache/ip/*.txt`, and `adguard/work/data/filters/*.txt` (mode
1611 `0444`). Run it once after setup or after pulling an old repo; `toggle.sh` and `sync-rules.py` will then
1612 work without permission errors.
1613
1614 **Context:** The nullexit project has a strict project-wide policy of `no chmod` from scripts (see README 6).
1615 This rule exists because every prior `chmod 0444` (post-write tamper-proof lock) and `chmod 000` (post-
1616 toggle `.env` lock) call from `scripts/sync-rules.py` / `toggle.sh` could leave a file permanently
1617 unreadable on disk if the script exited early, was killed mid-flight, or simply set a restrictive mode
1618 without ever being re-flipped. We've stripped every `chmod` from the codebase. But files **written by older
1619 versions of those scripts** are still on disk in those restrictive modes (`0444` for `compiled_rules.txt`,
1620 `ip_blocklist.ipset`, `data/filters/<id>.txt`; `000` for `.env` after end-of-toggle lock). Re-running `
1621 toggle.sh` or `scripts/sync-rules.py` against those leftovers hits `PermissionError: [Errno 13] Permission
1622 denied` immediately.
1623
1624 **Why the stale permissions also block `mv`/`rm` of the parent folder:** The restrictive modes (`000` on `.env
1625`, `0444` on output files) don't just block writes they prevent the kernel from unlinking the file's
1626 dentry during a `mv` or `rm` of the **containing directory** (`adguard/work/`). macOS enforces this at the
1627 VFS layer: you own the directory, but you can't delete a directory that contains entries you can't write to
1628 . This made the entire `adguard/work/` tree unmovable/undeletable until the locked inodes inside it were
1629 replaced. `unlock-files.sh` resolves this permanently by swapping each locked inode for a fresh `0644` one
1630 via atomic rename (`cp` + `mv`), which operates on the directory entry rather than the file contents no `
1631 chmod` called, no permission bypass needed.

```

```

1596 **Problem:** You (the owner) cannot `cat`, `cp`, `rm`, `>` redirect, or any normal POSIX write against a `mode
=000` file even though you own it the basic mode bits block ALL access for owner, group, and other. The
script does not call `chmod` and you want a permanent fix that never relies on a chmod from any future run
either.
1597
1598 **Solution:** Replace the locked inode with a fresh readable one. `mv` is atomic; mode bits live in the inode,
not the directory entry. Two flavors cover every case we hit on a real machine:
1599
1600 **Flavor A locked file is mode `000` (owner can't even READ it):**
1601 NOPASSWD sudoers (`/etc/sudoers.d/nullexit`) already includes `/usr/bin/python3`. So we read via `sudo python3`
(which has the DAC override root has) and pipe the bytes around as base64 in user space, where the
redirect happens with the user's normal umask = `0644`:
1602 ```bash
1603 sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).
decode())" > /tmp/snap.b64
1604 base64 -d < /tmp/snap.b64 > <FILE>.new
1605 mv -f <FILE>.new <FILE>
1606 ls -la <FILE> # mode is now 0644
1607 ```
1608 The old `mode=000` inode is unlinked by `mv`; the new `mode=0644` inode (created by base64-decode via the
calling user's umask) takes the path. No chmod ever called. The parent directory just needs to be writable
by you (it always is, since you own it).
1609
1610 **Flavor B locked file is mode `0444` (you can READ but cannot WRITE; `scripts/sync-rules.py` needs to re-
write):**
1611 You own the file and can read it, you just cannot truncate/overwrite it. The simplest path is to rename it
aside and let the writer create a fresh inode from scratch (which gets the umask-default `0644`):
1612 ```bash
1613 mv <FILE> /tmp/old.<FILE>.$$
1614 python3 scripts/sync-rules.py # now writes <FILE> cleanly at mode 0644
1615 ```
1616 Or apply Flavor A if you'd prefer to preserve the contents while resetting the mode.
1617
1618 **Why this is the canonical recovery, not a hack:**
1619 - Mode bits live in the inode, not the path. `mv -f` replaces the path's inode reference; the old locked inode
drops out of the directory entirely (dentry eviction) and loses its last link, so the kernel reclaims it.
The new inode (whatever it is: a freshly-written one, or a base64-decoded copy) gets the path.
1620 - The only prerequisite to apply Flavor B's `mv` is: you own the parent directory (so you can unlink entries
from it) AND you have a NOPASSWD-capable binary that can read the byte stream if mode prohibits user read.
1621 - This pattern generalizes. Any time a script ends up producing a "locked file that the next run can't update"
situation, the same trick applies without a single chmod anywhere.
1622
1623 **Empirical verification (user machine, July 1, 2026):**
1624 - `.env` was `mode=000` (owner $USER, i.e. you): Flavor A unblocked it in 3 commands, no chmod.
1625 - Before: `----- 1 $USER staff 899 Jun 28 07:07 .env`
1626 - After: `-rw-r--r--@ 1 $USER staff 899 Jul 1 20:39 .env`
1627 - `docker compose config` immediately picked up `TS_AUTHKEY` + `WIREGUARD_PRIVATE_KEY` substitution.
1628 - `compiled_rules.txt`, `ip_blocklist.ipset`, `data/filters/1782645604.txt` were `mode=0444`: Flavor B (`mv`-
aside) unblocked all three.
1629 - `scripts/sync-rules.py` then ran clean: 279587 block rules + 171 allow rules + 16710 IP entries compiled; new
files came out at `0644`.
1630 - `toggle.sh` then went end-to-end in 48s (warp / routing-fix / adguardhome / tailscale all healthy, gateway
mesh-joined, DNS hijack verified single-server).
1631
1632 **When this trick is appropriate vs. inappropriate:**
1633 - You own the parent directory and the file (typical desktop scenario).
1634 - The lock is a leftover from a removed `chmod` call that you want off disk forever.
1635 - NOPASSWD sudo includes at least one interpreter (`python3` on this system) that can be used to read mode-0
files. If it does not, add it first: `**Reusable cookbook:**
1641
1642 ```bash
1643 # Quick diagnostic: figure out which flavor you need.
1644 WHOAMI=$(whoami)
1645 ls -la <FILE>
1646 stat -f 'mode=%Sp owner=%Su group=%Sg' <FILE>
1647 # mode=----- or mode=----r----- : you cannot even read it Flavor A.
1648 # mode=r--r--r-- but writer fails : you are blocked from write but can read Flavor B.
1649 # mode=rw-r--r-- : not actually locked at all; unrelated write failure.
1650
1651 # Flavor A (mode=000):
1652 sudo -n /usr/bin/python3 -c "import base64,sys; sys.stdout.write(base64.b64encode(open('<FILE>','rb').read()).
decode())" > /tmp/snap.b64
1653 base64 -d < /tmp/snap.b64 > <FILE>.new && mv -f <FILE>.new <FILE> && rm -f /tmp/snap.b64

```

```

1654
1655 # Flavor B (mode=0444, file is readable):
1656 mv <FILE> /tmp/old.<FILE>.$(date +%s)
1657 # (let the original writer recreate <FILE>; new file lands at mode 0644)
1658
1659 # Verify after either flavor:
1660 ls -la <FILE> # mode should be 0644
1661 <your normal validation command>
1662 ```
1663
1664 ### 15.10 Tor Container
1665
1666 ##### 15.10.1 July 11, 2026: Tor Container Hardening & obfs4proxy Restoration
1667 **Symptom:**
1668 1. The Tor container entered a fatal crash loop upon boot, throwing `exec: "obfs4proxy": executable file not
    found in $PATH` when `TOR_USE_BRIDGES=true`.
1669 2. When the user pasted Tor bridge lines containing comments (`#`) or blank lines into `tor-bridges.txt`, the
    container failed to validate the file properly and crashed.
1670
1671 **Root Cause:**
1672 - **Bug 1 (Missing Pluggable Transports):** The Tor Dockerfile was built on `debian:bullseye-slim`, which had
    silently dropped or failed to resolve the `obfs4proxy` package in its latest apt repositories.
1673 - **Bug 2 (Fragile Validation):** The `entrypoint.sh` bridge validation check merely checked `[ -s tor-bridges.
    txt ]` (file size > 0). It did not verify if the file actually contained valid, uncommented bridge lines.
    Passing empty lines or comments tricked the validation script into appending invalid `Bridge` directives to
    the `torrc`, causing the Tor daemon to instantly exit.
1674
1675 **Resolution:**
1676 - **Upgraded Base Image:** Shifted the Tor Dockerfile base image to `debian:bookworm-slim`, restoring access to
    the `obfs4proxy` pluggable transport package.
1677 - **Robust Bridge Validation:** Replaced the fragile `[ -s ]` check with a strict `grep -vE '\s*(#|$)'\`
    command that strips comments and empty lines. If no valid lines remain, the container safely falls back to
    standard public guard relays instead of crashing, ensuring Tor remains available even with a misconfigured
    bridge file.
1678 - **Image Pinning:** Explicitly pinned `image: nullexit-tor:v1.0.0` in `docker-compose.yml` to prevent
    arbitrary local builds from drifting to `:latest`.
1679
1680 ##### 15.10.2 July 12, 2026: Tor ControlPort Dynamic Password Authentication & User Config
1681 **Symptom:** Querying the Tor Control Port (e.g., via the `/sweep` script or `check_circuits.py`) returned
    empty responses or connection errors. Logs showed that Tor automatically closed its ControlPort:
1682 `You have a ControlPort set to accept unauthenticated connections from a non-local address. ... That's so bad
    that I'm closing your ControlPort for you.`
1683
1684 **Root Cause:** Docker port mapping (especially under Colima on macOS) requires container sockets to bind to
    `0.0.0.0` to be reachable from the host. However, when Tor's ControlPort is bound to `0.0.0.0` (non-local)
    with no authentication, Tor closes it by default as a safety precaution. Sourcing SOCKS5/ControlPort to
    `127.0.0.1` inside the container was not an option as it broke the Docker bridge routing path.
1685
1686 **Resolution:**
1687 - **Dynamic Password Authentication:** Updated the Tor container's `entrypoint.sh` to read `TOR_PASSWORD` and
    dynamically generate its HMAC hash via `tor --hash-password`, adding `HashedControlPassword <hash>` to `
    torrc`. This secures the ControlPort on `0.0.0.0` so Tor keeps it open.
1688 - **Unified Credentials in `.env`:** Exposed `ADGUARD_USER=admin` and `ADGUARD_PASSWORD=nullexit` in `.env` to
    serve as the unified gateway credentials. Sourced `TOR_PASSWORD` directly from `ADGUARD_PASSWORD` in `
    docker-compose.yml`.
1689 - **Client Authentication:** Updated `scripts/check_circuits.py` to fetch `TOR_PASSWORD` (defaulting to `
    ADGUARD_PASSWORD` or `nullexit`) and authenticate with the ControlPort using `AUTHENTICATE "<password>"` to
    regain access.
1690
1691 > See also 11.3 (Tor Container Kill-Switch Inheritance & Fail-Closed Egress Verification) for the design-level
    fail-closed guarantee.
1692
1693 ### 15.11 macOS Platform Quirks
1694
1695 ##### 15.11.1 macOS Application Firewall Silently Blocking AirDrop & Continuity
1696 **Problem:** Users of complex network extensions (Tailscale, LuLu, Docker, etc.) often run into an issue where
    AirDrop, AirPlay, and Universal Clipboard silently fail, even with the gateway inactive and Wi-Fi/Bluetooth
    correctly configured. The internal Apple Wireless Direct Link (`awdl0`) interface appears healthy, but
    connections never establish.
1697
1698 **Root Cause:** The built-in macOS Application Firewall has a specific list of explicitly blocked applications.
    Apple's own core networking services (e.g., `/usr/libexec/sharingd`, `/usr/libexec/rapportd`, `/usr/
    libexec/avconferenced`) can be silently added to the "Block incoming connections" list either through
    accidental "Deny" clicks on permission prompts, execution of aggressive privacy-hardening scripts, or a
    known macOS bug that retains strict blocks even after toggling the "Block all incoming connections" setting
    off. Since these are background daemons, macOS provides no visible error.
1699
1700 **Fix:** The firewall rules must be cleared via the terminal (or System Settings > Network > Firewall > Options
    ):
1701 ```bash
1702 sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/sharingd

```

```

1703 sudo /usr/libexec/ApplicationFirewall/socketfilterfw --unblockapp /usr/libexec/rappor
1704 tdd
1705 Once unblocked, `sharingd` immediately resumes broadcasting mDNS/Bonjour discovery packets, and AirDrop
    restores instantly without requiring a reboot.
1706
1707 ##### 15.11.2 Standalone Tailscale Daemon Freeze after macOS Sleep/Wake
1708 **Problem:** When the macOS host goes to sleep and wakes up, the network interfaces and routing tables are
    rebuilt. The standalone `tailscaled` daemon (installed via Homebrew) occasionally fails to handle this
    transition and becomes completely frozen/unresponsive. In this state, any command like `tailscale down` or
    `tailscale status` hangs indefinitely, and all DNS requests sent to the local Tailscale resolver
    (`100.100.21.8`) time out, causing a total internet blackout on the host.
1709
1710 **Fix:** Enhanced the Tailscale verification and teardown logic in `toggle.sh` and `recover.sh`:
1711 1. **Unresponsive Daemon Detection:** Instead of assuming the daemon is healthy just because the Unix socket `
    var/run/tailscaled.socket` exists, the scripts now actively run a status check with a 5-second timeout (`
    run_with_timeout 5 tailscale status`). If the check times out (exit code 143), the daemon is classified as
    unresponsive.
1712 2. **Auto-Recovery Loop:** If the daemon is unresponsive, the script now automatically attempts to restart the
    service using `brew services restart tailscale` (falling back to `sudo -n brew services restart tailscale`)
    .
1713 3. **Graceful Retry:** Once restarted, the script pauses for 3 seconds for the daemon to initialize before
    proceeding with connection or teardown commands.
1714
1715 ##### 15.11.3 macOS Sharing Services (AirDrop/AirPlay) Freeze on Network Transitions
1716 **Problem:** When the gateway is turned on or off, the scripts perform network configuration cleanups which
    include flushing the routing tables (`route -n flush`), restarting the `mDNSResponder` daemon (`killall -
    HUP mDNSResponder`), and power-cycling the Wi-Fi interface (`setairportpower` or `ifconfig en0 down/up`).
    While AirDrop BLE discovery continues to function, the actual file transfers stall at "Waiting..."
    indefinitely.
1717
1718 **Root Cause:** The rapid tear-down and rebuild of the local routing table and Wi-Fi interface causes Apple's
    core sharing and connection daemons (`sharingd` and `rappor
1719 tdd`) to lose their socket bindings to the mDNS/
    Bonjour discovery interface. Instead of reconnecting gracefully, they enter a wedged state and try to route
    peer-to-peer (AWDL) IPv6/TCP transfer traffic into the default gateway (the Tailscale interface `utun`)
    instead of the direct wireless link, causing the TLS verification handshake to time out.
1720
1721 **Fix:** Integrated an automatic sharing services reset into both `toggle.sh` and `recover.sh`:
1722 1. The script now automatically runs `sudo -n killall sharingd rappor
1723 tdd` at the end of both the **START** path
    and the **STOP** path (as well as inside `recover.sh`).
1724 2. Killing these daemons forces macOS to immediately relaunch them, binding them fresh to the newly initialized
    network interfaces and routing tables, allowing AirDrop to work seamlessly while the gateway is active.
1725
1726 ##### 15.11.4 Chrome Remote Desktop Connection Failures (Tailscale on Remote Device)
1727 **Context:** When attempting to access a remote machine via Chrome Remote Desktop (CRD), the connection fails
    or hangs if Tailscale is active **on the remote device**. The connection works fine even if Tailscale (and
    the exit node) is active **on the local client/viewer device**, but only if Tailscale is disabled on the
    remote machine. *(This is distinct from the CRD-through-WARP latency limitation in 14.3.)*
1728
1729 **Why:** Having Tailscale active on the remote machine creates virtual network interfaces and DNS/routing
    modifications that conflict with Chrome Remote Desktop's WebRTC bindings and signaling traffic.
1730
1731 **Workaround:**
1732 - **Use Tailscale Native RDP/RustDesk (Recommended):** Connect directly to the remote device's Tailscale IP
    (`100.X.Y.Z`) via an RDP or RustDesk client. This is faster and avoids third-party relay conflicts.
1733 - **Disable Tailscale on the Remote Device:** If you must use CRD, disable Tailscale on the remote machine
    before connecting.
1734
1735 ##### 15.11.5 tailscaled Daemon Broken State (brew services)
1736 **Symptom:** `brew services list` shows `started` but `tailscale status` shows "stopped".
1737
1738 **Fix:** `sudo brew services restart tailscale`. Toggle script now detects and auto-restarts.
1739
1740 ##### 15.11.6 Apple Silently Removed the `airport` Binary in macOS Sonoma 14.4 (March 2024)
1741 **What happened:** `watcher.sh`'s `detect_lan_p2p_mode()` was written to use the `airport` CLI tool at:
1742 ...
1743 /System/Library/PrivateFrameworks/Apple80211.framework/Versions/Current/Resources/airport
1744 ...
1745 This path returned `exit 127: no such file or directory`. The function silently fell back to `reason="no-wifi"`
    even while the Mac was actively connected to WPA2 Enterprise Wi-Fi.
1746
1747 **Why Apple removed it:** `airport` was never a real public binary it lived inside a **private framework** and
    was always technically unsupported. Apple deprecated it quietly years ago and finally removed it in **
    macOS Sonoma 14.4 (March 2024)**. The removal is part of a broader privacy clampdown on Wi-Fi metadata:
1748 - Starting in **macOS 14.5**, BSSIDs and other Wi-Fi association details are **redacted** (`<redacted>`) in
    most tools, including the official replacement `wdutil`
1749 - Apple's official recommendation is to use the **CoreWLAN framework** (Swift/Objective-C) with proper
    entitlements useless for a bash script
1750 - The official CLI replacement `wdutil` requires `sudo` for most useful output also useless for a background
    LaunchAgent
1751
1752 **Fix:** Replaced `airport -I | awk '/link auth/'` with:

```



```

1751 ```bash
1752 system_profiler SPAirPortDataType | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF
1753 ; exit}'
1754 ```
1755 This returns human-readable strings like `WPA2 Enterprise`, `WPA2 Personal`, or `None` for the currently
1756 associated network without `sudo`, and without any removed private binary.
1757
1758 **Confirmed working output on this machine:**
1759 ```
1760 P2P detect: security='Enterprise' allow=false (reason: 802.1x-enterprise)
1761 ```
1762
1763 > ** Risk: `system_profiler` may not survive forever either.** `system_profiler SPAirPortDataType` still works
1764 as of macOS Sequoia (15.x) without sudo and without redaction. However, given Apple's trajectory on Wi-Fi
1765 privacy, this could be locked down in a future release. If `security` comes back empty on a future macOS
1766 version while the Mac is actively on Wi-Fi, the function will silently fall back to `allow=false` (safe
1767 default), but the detection will be broken again. The long-term fix would be a small Swift helper using
1768 CoreWLAN that we compile once and bundle in the repo. *(This forward-looking caveat is also tracked as an
1769 observation in 16.)*
1770
1771 ##### 15.11.7 Changing macOS Local Hostname
1772 **Context:** Users may wish to change their Mac's computer name/hostname in macOS System Settings.
1773 **Effect:** Changing the Mac's hostname will not break Nullexit (which relies on internal routing/localhost
1774 ). However, Tailscale will automatically detect this change and update the machine name on the Tailnet.
1775 - The underlying Tailscale `100.x.x.x` IPv4 address will remain exactly the same.
1776 - The MagicDNS name will change to match the new hostname. Users must remember to update their SFTP/SSH
1777 clients to use the new MagicDNS address.
1778
1779 ##### 15.11.8 Shell-Language Landmines: macOS `/bin/bash` 3.2 + `set -e` (Field Notes)
1780
1781 This project's lifecycle scripts (`toggle.sh`, `recover.sh`, `common.sh`, ) run under `#!/bin/bash`, which on
1782 macOS is GNU bash 3.2.57 (2007). Apple has never shipped a newer bash because bash 4+ is GPLv3. (
1783 Homebrew's bash 5 lives at `/opt/homebrew/bin/bash` but is not the shebang target, and cannot be assumed
1784 on PATH a bare `bash --version` on a stock Mac reports 3.2.) Every script here also runs `set -e` (exit-on-
1785 -error) with an `ERR`/`INT`/`TERM`/`HUP` trap. That combination ancient bash + `set -e` + traps + `set -
1786 x` xtrace redirection is a minefield. The traps below were all hit and fixed during development;
1787 document them so they are not re-discovered.
1788
1789 **A. `set -e` + `if then else fi` can abort at the `else` (bash 3.2 heisenbug).**
1790 The nastiest one. A construct as innocent as:
1791 ```bash
1792 if [ -f "$MARKER_FILE" ]; then
1793     X="yes"
1794 else
1795     X="no"
1796 fi
1797 ```
1798 aborted `toggle.sh` at init under `set -e` whenever the file was absent the false status of the `[ -f ]`
1799 test leaked through the compound and `set -e` killed the script (`EXIT: code=1 phase=init`), before any
1800 real work ran. Two things made this vicious: (1) it is intermittent/context-dependent it did not
1801 reproduce in isolated `bash -c` snippets of the identical block; it only manifested inside the full
1802 script with the `set -x` xtrace fd-redirect (`exec 2> >(tee)`) active. We only proved causation by 
1803 bisecting the live script (replace the block with a trivial assignment the script sailed past init). (2)
1804 Static tools (`bash -n`, `shellcheck`) cannot see it it is a runtime interaction. Safe pattern:
1805 default-assign, then an `if` with no `else` (an `if` whose false condition has no else-branch yields
1806 status 0):
1807 ```bash
1808 X="no"
1809 if [ -f "$MARKER_FILE" ]; then X="yes"; fi
1810 ```
1811
1812 **Confirmed root cause (2026-07-13, reproduced head-to-head).** The trigger is specifically set -x plus a
1813 `PS4` that contains a command substitution here the `DEBUG_TRACE` prompt `PS4='+ [$(date)] '`. Forcing
1814 the false-condition/else path (`gateway_state` `stopped`, so `if [ "$(gateway_state)" = "running" ]` takes
1815 the else/START branch) on stock `/bin/bash` 3.2.57: with the default `PS4` it survives (3/3, exit 0);
1816 with toggle's `${(date)` `PS4` the `ERR` trap fires on the else branch (`rc=1` from the false condition,
1817 3/3, exit 42). The `${(date)` runs a subshell on every traced command; while tracing the if-condition
1818 it perturbs `${(date)` so the condition's non-zero status leaks past the normal if-exemption and trips the `ERR`
1819 trap. This is a true heisenbug enabling tracing (`DEBUG_TRACE=true`) is what causes the abort,
1820 which is exactly why it's invisible with `DEBUG_TRACE=false` (the default) and why every isolated repro on
1821 the default `PS4` passes (this misled an entire debugging session into briefly "disproving" the landmine)
1822 . It aborted the gateway START decision (`toggle.sh:689`) under `DEBUG_TRACE=true`; with `DEBUG_TRACE=false`
1823 the gateway starts cleanly (verified: 192 s, all five containers healthy). Fixed (2026-07-14).
1824 DEBUG_TRACE's `PS4` now uses the subshell-free `${SECONDS}` (elapsed seconds since the shell started)
1825 instead of `${(date)` removing the per-traced-command subshell that perturbed `${(date)`. bash 3.2
1826 has no subshell-free `${wall-clock}` token for `PS4` (`${EPOCHSECONDS}` is 5.0+, `${printf '%()T'}` is 4.2+), so
1827 the per-line stamp is now elapsed-seconds the absolute start time is still logged once in the DEBUG_TRACE
1828 enabled header line, so wall-clock is recoverable. Bonus: this also drops a date fork on every
1829 traced line (previously thousands). Verified head-to-head on stock `/bin/bash` 3.2.57: toggle.sh --
1830 restart with DEBUG_TRACE=true now completes code=0 with tracing active straight through the exact
1831 decision that aborted before the trace shows gateway_state stopped (`common.sh:229`) feeding the false
1832 if [ "$(gateway_state)" = "running" ], then the else branch reaching Action: STARTUP GATEWAY (1436

```



traced lines, no `ERR` trap firing, all five containers healthy). `DEBUG\_TRACE` is now safe to use on the STOP/START decision path.

More generally, any command whose "failure" is normal control flow (`grep -q`, `[ ]`, `diff`, `kill -0`) that ends up as the *last* command of a function/branch/script under `set -e` is a landmine; guard it (`|| true`) or restructure so a non-zero status can't be the tail value.

**B.** `BASH\_XTRACEFD` and the `exec {var}>>file` dynamic-fd form are bash 4.1+ (hard-fail on 3.2).  
Our `DEBUG\_TRACE` feature wants xtrace (`set -x`) sent to `output.log` while keeping the terminal clean the clean way is `exec {BASH\_XTRACEFD}>>"\$LOG\_FILE"`. On 3.2 this is a *syntax/runtime error*: `exec: {BASH\_XTRACEFD}: not found` (dynamic-fd allocation `{var}>>` didn't exist until 4.1, and `BASH\_XTRACEFD` itself is 4.1+; on 3.2 it's just an ordinary variable that `set -x` ignores). Left unguarded it would abort the script at startup i.e. the debug feature would break the very thing it's meant to observe. **Safe pattern**: *branch* on `\${BASH\_VERSINFO[0]}.\${BASH\_VERSINFO[1]}` on 4.1 use `exec 9>>"\$LOG\_FILE"; export BASH\_XTRACEFD=9`; on 3.2 fall back to `exec 2> >(tee -a "\$LOG\_FILE" >&2)` (xtrace goes to stderr, which we duplicate to the log; the terminal also shows it, which is acceptable). Never use the `exec {var}>>` form unless you *know* you're on 4.1.

**C.** `trap ERR` + `\${LINENO}` misattributes the failing line on 3.2.  
When a command fails under `set -e`, our trap runs `cleanup\_handler ERR \$LINENO`. On bash 3.2, `\${LINENO}` inside an `ERR` trap frequently reports the line of the *enclosing* `fi`/closing brace of a compound statement, not the true failing line. During the STOP/START-decision bug this produced `Script failed at line 1202` where 1202 is a bare *fi* actively misleading. **Mitigation**: don't trust the trap's line number as gospel on macOS; corroborate with the `EXEC:/EXIT:` lifecycle breadcrumbs and the `DEBUG\_TRACE` xtrace (which prints *file:line:function* via a custom `PS4`), which pin the *actual* last-executed command.

**D.** `set -e` is not inherited by command substitutions / subshells the way people expect, and is suppressed inside *if* conditions.  
This cuts both ways and is a frequent source of "why didn't it fail?" and "why *did* it fail?": a non-zero command used as an *if* condition (or left/right of *if*) is *exempt* from `set -e` (by design), but the moment that same call is a bare statement it aborts. *is\_gateway\_active()* relies on this (it's always called as *if is\_gateway\_active*; then `), which is precisely why moving the decision to a marker file instead of a probe that must be *if*-guarded to be safe is more robust (see 15.12.19). Also note *local X=\$(cmd)* masks *cmd*'s exit status (the *local* builtin's status wins) a *shellcheck* SC2155 we see repeatedly; harmless when we don't check ` \$?`, but a real footgun if you do.

**E.** Process substitution (*>>()*, *<<()*) is a bashism, and its lifecycle is loose.  
Our 3.2 xtrace fallback (*exec 2> >(tee -a "\$LOG\_FILE" >&2)*) uses process substitution *fine* under bash, but (1) it is *not* POSIX *sh*, so these scripts must never be invoked as *sh script.sh*; and (2) the background *tee* it spawns is reaped asynchronously, so the very last few lines written just before exit can occasionally be truncated in the log. Acceptable for a debug trace; do not rely on it for critical last-gasp logging (use a direct *>> "\$LOG\_FILE"* append for must-capture lines, as the *EXIT* breadcrumb does).

**F.** GNU-vs-BSD userland divergence (not bash itself, but same class of pain).  
macOS ships BSD variants of the core tools, so idioms differ from Linux: *sed -i ''* (BSD requires the empty backup-suffix arg) vs *sed -i* (GNU); *stat -f%z* (BSD) vs *stat -c%s* (GNU); *date* format flags; *jot* (BSD) vs *shuf* (GNU) for random port generation; *route* / *networksetup* / *dscacheutil* (macOS-only) vs *ip* / *resolvectl* (Linux). This is *why* the codebase branches on *[[ "\$OSTYPE" == "darwin\*" ]]* everywhere *toggle.sh* and *recover.sh* are each single cross-platform scripts that dispatch on *\$OSTYPE* (an early Linux block runs and exits before the macOS body). *timeout(1)* doesn't exist on stock macOS at all hence the pure-bash *run\_with\_timeout* in *common.sh*.

**Working rules distilled from the above**:  
- Assume *bash 3.2* semantics for anything under *#!/bin/bash* on macOS; gate any 4.0+ feature (*BASH\_XTRACEFD*, dynamic *{var}* fds, associative arrays, `\${var}` case-conversion, *mapfile* / *readarray*, negative array indices) behind a *BASH\_VERSINFO* check or avoid it.  
- Under *set -e*, never let a "benign non-zero" command ([ ], *grep -q*, *kill -0*, *diff*) be the tail statement of a branch/function/script; use *no-else if*'s, *|| true*, or explicit  *\$?*  handling.  
- Prefer *persisted state* (marker files) over *live probes* for control-flow decisions a probe forces an *if*-guard and can hang/abort; a file read cannot (see 15.12.19).  
- *bash -n* and *shellcheck* catch syntax and many smells but *cannot* catch *set -e* / trap runtime interactions a *live* run is mandatory to validate lifecycle-script changes; the *DEBUG\_TRACE* + breadcrumb instrumentation exists precisely to make those live runs diagnosable.  
- Never run these scripts as *sh* they are bashisms end to end.

### 15.12 Misc Resolved

#### 15.12.1 Gluetun Resets nftables FORWARD Rules

**Symptom**: FORWARD RELATED, ESTABLISHED rule disappears after Gluetun health check.

**Fix**: *scripts/routing-fix.sh* re-adds to both iptables backends every 30 seconds. Legacy backend (policy ACCEPT) acts as safety net.

#### 15.12.2 Native SSH & SFTP Security (Zero Local Attack Surface)

**Context**: macOS "Remote Login" (*sshd*) and "File Sharing" (*smbd*) open ports on the local network (e.g. *172.x.x.x*), exposing the host to brute-force attacks on public Wi-Fi.

**Implementation**: The script strictly enforces *tailscale up --ssh=true* whenever the mesh state is reset. This empowers the user to completely disable native macOS Remote Login and File Sharing. This provides an incredible zero-trust security advantage: even if an attacker steals your Mac username and password, they *cannot* access your files remotely. Disabling the native services completely closes the listening ports

on the local Wi-Fi interface (`172.x.x.x`), rendering the Mac inaccessible to local network logins. Meanwhile, Tailscale intercepts port 22 traffic exclusively over the encrypted mesh and authenticates cryptographically based on your Tailscale SSO identity. Stolen Mac passwords are mathematically useless without first bypassing your Tailscale Two-Factor Authentication and registering a device onto the mesh.

1824

1825 **\*\*Cross-Platform File Exchange:\*\*** To easily browse and exchange files securely, simply use an SFTP client and connect to your Mac's MagicDNS name on port 22. Recommended clients: **\*\*WinSCP\*\*** or **\*\*FileZilla\*\*** (Windows), **\*\*Cyberduck\*\*** (macOS), and **\*\*FE File Explorer\*\*** or **\*\*Solid Explorer\*\*** (iOS/Android).

1826

1827 **#### 15.12.3 Logging Architecture**

1828 For debugging, logs are strictly segmented based on the component's lifecycle:

1829 - **\*\*`output.log` (Host-side):\*\*** Contains all standard error (`stderr`) and verbose output from the host scripts (`toggle.sh`, `recover.sh`, `setup.sh`). Since the `rule-compiler` container is ephemeral and deleted after running, its logs (and any Python errors from `scripts/sync-rules.py`) are extracted and appended to this file before deletion.

1830 - **\*\*Log Rotation Policy:\*\*** On startup, `toggle.sh` checks if `output.log` exceeds **\*\*1GB\*\*** (`1,073,741,824` bytes). If it does, it performs a metadata-only rename (`mv output.log output.log.old`), preserving history while resetting the active log to 0 bytes. This rename-based rotation avoids the heavy disk I/O and CPU overhead of rewriting a massive file line-by-line. The threshold is intentionally large because with `DEBUG\_TRACE=true` the log accumulates a full command-by-command execution trace (effectively the framework's entire compilation/warning history), which is valuable to retain for post-mortems; a smaller cap would rotate that history away too quickly.

1831 - **\*\*Egress Probe Logging:\*\*** The background host-egress leak probe checks if stdout is a TTY (`[ -t 1 ]`) and will only print live status updates (using carriage return `\r`) or duplicate console warnings in interactive mode. It automatically detects macOS/BSD vs. Linux environments to dynamically construct timestamps with date and time (omitting milliseconds on macOS to prevent high CPU utilization from spawning sub-second processes).

1832 - **\*\*`docker logs <container>` (Guest-side):\*\*** The persistent containers (`warp`, `tailscale`, `routing-fix`, `adguardhome`) use the Docker `json-file` logging driver with a strict `max-size` (1m-10m) to prevent VM disk exhaustion. Use standard `docker logs` to view them.

1833

1834 **#### 15.12.4 Pre-Flight Check 1: Wrong `tailscale ping` Flag**

1835 **\*\*Problem:\*\*** Check 1 of the pre-flight connectivity checks used `--c 1` (double dash) but `tailscale ping` accepts `-c 1` (single dash). The invalid flag caused the command to exit with code 1, marking the check as FAIL even though the gateway was reachable.

1836

1837 **\*\*Fix:\*\*** Changed `--c 1` to `-c 1`.

1838

1839 **#### 15.12.5 Pre-Flight Check 1: DERP Relay Exit Code**

1840 **\*\*Problem:\*\*** Even with the correct `-c 1` flag, `tailscale ping` exits with code 1 when the connection goes through a DERP relay (`direct connection not established`) even though a pong was successfully received (28ms via DERP(tor)). The check treated relayed connections as failures.

1841

1842 **\*\*Fix:\*\*** Changed the check from relying on exit code to piping stdout through `grep -q "pong"`. This accepts relayed connections as valid (they work fine for the exit node use case), while still failing on true timeouts or unreachable peers.

1843

1844 **#### 15.12.6 IPv6 Exit-Node Leak**

1845 **\*\*Problem:\*\*** Tailscale supports IPv6, but WARP/gluetun is IPv4-only. IPv6 packets forwarded through the exit node would leak through the Docker bridge unencrypted.

1846

1847 **\*\*Fix:\*\*** Added `ip6tables -A FORWARD -i tailscale0 -j DROP` to `post-rules.txt`, blocking all IPv6 forwarded traffic from the Tailscale interface. \*(This is the fix referenced by the IPv6-leak scenario in 15.6.1.)\*

1848

1849 **#### 15.12.7 SOCKS5 Proxy Threading Race Condition**

1850 **\*\*Problem:\*\*** The `forward` function in `scripts/socks5-proxy.py` called `select.select([src, dst], [], [])` which raised `ValueError: file descriptor cannot be a negative integer (-1)` when one thread closed a socket while the other thread was selecting on it.

1851

1852 **\*\*Fix:\*\*** Added `ValueError` exception handling around `select.select()` and `recv()` calls. When a file descriptor becomes negative, the thread breaks out of the forwarding loop cleanly instead of crashing.

1853

1854 **#### 15.12.8 Docker Healthcheck: Multi-line Python in CMD Array**

1855 **\*\*Problem:\*\*** The socks-proxy healthcheck used a `CMD` array with a multi-line Python string. YAML collapsed the newlines, causing the `#` comment to comment out the rest of the code and the `assert` statement to be mangled into `as`.

1856

1857 **\*\*Fix:\*\*** Changed from `["CMD", "python3", "-c", "..."]` to `["CMD-SHELL", "python3 -c '...']` with a single-line Python expression. Uses a real SOCKS5 handshake (sends `[5, 1, 0]`, expects `[5, 0]`) instead of a simple TCP port check.

1858

1859 **#### 15.12.9 Graceful Shutdowns Leave DNS Hijacked**

1860 **\*\*Problem:\*\*** The gateway overrides the macOS host's DNS settings (via `networksetup`) to route queries to the AdGuard container. Because macOS persists `networksetup` changes to disk, shutting down the Mac while the gateway is running causes the Mac to boot up later with hijacked DNS. Since the Colima VM and Docker containers don't auto-start on boot, the user would have no internet access until they manually run the toggle script.

1861

1862 **\*\*Root Cause:\*\*** The `toggle.sh` script exits after the gateway starts. There was no background daemon listening for macOS system shutdown signals to revert the network settings.

1863

1864 **\*\*Fix:\*\*** Wrapped the background `caffeinate` process (used to prevent system sleep) in a bash `trap` command. The trap listens for `SIGTERM`, `SIGINT`, and `SIGHUP`. When the user initiates a graceful shutdown, macOS sends `SIGTERM` to the background trap, which instantly executes `recover.sh` to flush the host DNS back to `1.1.1.1` right before the machine powers off.

1865 **\*Note:\*\*** We explicitly use `recover.sh` instead of `toggle.sh` because during a shutdown, the OS limits process cleanup time and is already independently sending termination signals to Docker and Colima. Trying to run `docker compose down` inside a shutdown trap creates a race condition that can hang the script until macOS forcefully `SIGKILL`s it. `recover.sh` abandons container management and surgically flushes the macOS network settings in milliseconds, guaranteeing completion before the OS pulls the plug.

1866

1867 ##### 15.12.10 Linux Native Wi-Fi Bouncing

1868 **\*\*Problem:\*\*** The generated `scripts/recover-linux.sh` incorrectly retained the macOS `networksetup` commands, which would crash and fail to bounce Wi-Fi on Linux hosts.

1869

1870 **\*\*Fix:\*\*** Swapped the macOS logic for native Linux commands. The script now attempts `nmcli radio wifi off/on` first, and gracefully falls back to `ip link set wlan0 down/up` if NetworkManager is unavailable.

1871

1872 ##### 15.12.11 TS\_AUTH\_ONCE Cloud Footgun & Ephemeral Keys

1873 **\*\*Decision:\*\*** The compose file uses `TS\_AUTH\_ONCE=true` to prevent authentication loops. However, on headless cloud VPS deployments, if a standard auth key expires (usually 90 days), the container will silently drop off the mesh. We documented this trade-off explicitly in `docker-compose.yml` and recommended the use of **\*\*Tailscale Ephemeral Keys\*\*** for cloud deployments, which automatically clean themselves up and bypass this issue.

1874

1875 ##### 15.12.12 Hardcoded AdGuard Credentials

1876 **\*\*Decision:\*\*** AdGuard is deployed with the hardcoded credentials `admin / nullexit`. This is perfectly safe behind a NAT on a local macOS laptop. However, if deployed on a cloud host with exposed ports, this becomes a severe security risk. We added a stark warning comment in `docker-compose.yml` to ensure cloud users change this before deployment.

1877

1878 ##### 15.12.13 Routing Fix: 30-second polling vs Netlink Socket

1879 **\*\*Decision:\*\*** Instead of building a complex Netlink socket listener (`ip monitor`) to react instantly to iptables and routing flushes from Gluetun, we retained the 5-second `sleep` loop in `scripts/routing-fix.sh`.

1880 **- \*\*Reasoning:\*\*** Building a netlink listener in pure bash on Alpine is incredibly brittle and complex (especially for iptables events). The 30-second polling loop is bulletproof. The only trade-off is a potential 1-4 second stutter if Gluetun reconnects, which was deemed an acceptable edge case for absolute reliability.

1881

1882 ##### 15.12.14 Cache Poisoning and URL Rot in scripts/sync-rules.py

1883 **\*\*Problem:\*\*** If a remote blocklist URL went permanently offline (404 Not Found), the Python `urllib` request would return the 404 HTML string. The script would aggressively regex-parse this HTML, find no domains, and overwrite the healthy local cache with a 0-domain file. This effectively disabled ad-blocking until the URL was fixed.

1884

1885 **\*\*Fix:\*\*** Implemented a defensive programming sanity check in `scripts/sync-rules.py`. Before overwriting the cache, the script now verifies that the compiled domain count is greater than 10 (and 1 for IP feeds). If it falls below this threshold (indicating a catastrophic 404 failure across multiple URLs), it aborts the cache overwrite, raises a `ValueError`, and keeps the previous day's healthy cache intact.

1886

1887 ##### 15.12.15 Gluetun nf\_tables Parser Crash (Bash Interpolation Failure)

1888 **\*\*Problem:\*\*** Attempting to make the TCP MSS dynamically configurable by using a standard bash variable inside `post-rules.txt` (`iptables ... --set-mss \${GATEWAY\_MSS:-1120}`) caused the Gluetun container to catastrophically crash during startup. Gluetun's internal parser executes `post-rules.txt` line by line without a shell, meaning the variable was never expanded. It literally attempted to inject the string `\${GATEWAY\_MSS:-1120}` into iptables, resulting in a fatal `bad value for option "--set-mss"` error.

1889

1890 **\*\*Fix:\*\*** Removed the bash variable from `post-rules.txt` and reverted it to a pure hardcoded integer. To retain dynamic `.env` configuration, we moved the interpolation logic to `toggle.sh`. Right before Docker starts, the host script parses `GATEWAY\_MSS` from the `.env` file and uses a cross-platform `sed` command to dynamically overwrite the integer directly inside `post-rules.txt`. This perfectly mimics manual configuration and completely bypasses Gluetun's strict parser.

1891

1892 ##### 15.12.16 AdGuard Filter Redundancy & Memory Optimization

1893 **\*\*Problem:\*\*** AdGuard Home does not perform cross-list deduplication in memory. When it loads its own native subscription filters (e.g., `AdGuard DNS filter`) alongside our massive `compiled\_rules.txt`, overlapping rules are loaded twice, wasting significant RAM in the Colima VM. The UI would show ~500k rules, representing the sum of all lists rather than unique domains.

1894

1895 **\*\*Fix:\*\*** Updated `scripts/sync-rules.py` to intelligently cross-reference AdGuard's configuration and deduplicate our list against it.

1896 1. The script now reads `adguard/conf/AdGuardHome.yaml` to dynamically fetch the exact URLs of any enabled native AdGuard filters.

1897 2. The parsing engine was upgraded to normalize basic AdGuard syntax (`||domain`) back into raw base domains during the build process.

1898 3. The script subtracts the native AdGuard domains from our custom blocklist *before* compiling, immediately purging ~83,000 completely redundant rules.

1899 4. The normalized base domains then pass through our subdomain optimizer, which squashes an additional ~50% of the remaining rules.

1900

This reduces the final compiled output from ~325k to ~281k rules on the `heavy` profile, saving ~45k rules from being loaded into memory twice and reducing the overall RAM footprint of the `adguardhome` container.

#### 15.12.17 Kernel-Level IP Blocklist (Threat Intelligence Firewall)

**\*\*Problem:\*\*** DNS sinkholing cannot stop malware or botnets that bypass DNS entirely by hardcoding direct IP addresses (a common technique for C2 communication). These connections pass straight through AdGuard unseen.

**\*\*Fix:\*\*** Added a second compilation pipeline to `scripts/sync-rules.py` that runs on every startup alongside the DNS pipeline.

1. The compiler concurrently fetches four curated threat-intelligence feeds: **\*\*Feodo Tracker\*\*** (abuse.ch botnet C2 IPs), **\*\*Spamhaus DROP\*\*** (IPs allocated to criminal organizations), **\*\*Emerging Threats\*\*** (Proofpoint active C2 and scanners), and **\*\*CINS\*\*** (active brute-force sources).
2. All entries are normalized via Python's `ipaddress` module. Any IP or CIDR that overlaps with RFC1918 private ranges, loopback, link-local, or the Tailscale CGNAT range (`100.64.0.0/10`) is stripped to prevent accidentally locking users out of their own LAN or mesh.
3. The cleaned list (16,721 unique IPs/CIDRs) is written as an `ipset restore` file using an **\*\*atomic swap pattern\*\***: `create\_new populate swap with live destroy\_new`. This guarantees the live `block\_malicious` ipset is never empty during a reload.
4. `scripts/routing-fix.sh` watches the file's mtime every 30 seconds. On change it runs `ipset restore` and idempotently re-injects `FORWARD DROP` rules for both `src` and `dst` into both iptables backends (nftables + legacy).
5. `docker-compose.yml` mounts `adguard/work/userfilters/` into `routing-fix` as `/userfilters:ro` and adds `rule-compiler: service\_completed\_successfully` to its `depends\_on`, eliminating the race condition where routing-fix could start before the file existed.

**\*\*Memory cost:\*\*** The entire 16,721-entry ipset costs only ~1.6 MiB of kernel memory negligible in the 600MB VM budget.

#### 15.12.18 Incident Post-Mortem: Discarded Docker Exec Errors in logs (July 10, 2026)

**\*\*Symptom:\*\*** When the `warp` container is stopped, dead, or failing to connect to its endpoints, the `toggle.sh` and `recover.sh` connectivity health checks fail, but no details are printed to `output.log`. The logs only show generic `FAIL` outputs, making it hard to see if the failure was a network timeout, Docker daemon error, or a missing binary.

**\*\*Root Cause:\*\*** Overly broad error silencing. The `docker compose exec` commands in the health checks (Check 3 in `toggle.sh` and the traffic check in `recover.sh`) both redirected stderr to `/dev/null` (`>/dev/null` and `2>/dev/null` respectively), completely throwing away any diagnostic error messages produced by Docker or `wget`.

**\*\*Fix (July 10, 2026):\*\*** Redirected stderr of the `docker compose exec` check commands to `output.log` (`2>> output.log`). This ensures that any command execution failures or connection timeout details are logged.

#### 15.12.19 Incident Post-Mortem: Toggle Decided STOP/START From a Live Probe, Aborting When Docker Was Down (July 13, 2026)

**\*\*Symptom:\*\*** Running `./toggle.sh` (or `--restart`) while the Colima VM / Docker daemon was **\*\*not\*\*** running caused the script to hang for ~15 seconds and then abort during the START phase without ever booting Colima. With `DEBUG\_TRACE` the lifecycle breadcrumb showed `toggle.sh EXIT: code=1 phase=start`, and the trace showed execution jumping straight from the START-branch entry to the `ERR` trap (`cleanup\_handler ERR`) with the START body never executing. Historically this silent failure occurred 57 in `output.log`. Spam-clicking the toggle during the resulting window (each retry rejected by the lock guard, then finally succeeding) is what produced the earlier "pile of stacked toggle processes" and Wi-Fi-bounce confusion.

**\*\*Root Cause:\*\*** `toggle.sh` decided whether to STOP or START by calling `is\_gateway\_active()` (`scripts/common.sh`), which **\*\*live-probes\*\*** the running system its first step is `run\_with\_timeout 15 docker compose ps --status running`. When `/var/run/docker.sock` is absent (Colima down), that call blocks for the full 15 s timeout and returns non-zero. Combined with the script's global `set -e` and `ERR` trap, a non-zero result evaluated at the `if is\_gateway\_active; then else fi` boundary aborted the whole script **\*\*before\*\*** the `else` (START) branch body ran precisely the branch whose job is to start Colima. Design-wise the deeper flaw is that a **\*\*disable\*\*** should not depend on whether the gateway is **\*\*currently\*\*** healthy; it should act on whether the gateway is **\*\*supposed\*\*** to be running (persisted intent). The correct signal already existed `MARKER\_FILE=/tmp/nullexit-gateway-active.marker`, written at the end of a successful START and cleared on STOP but `toggle.sh` only wrote/cleared it and never read it for this decision. **\*\*Note:\*\*** this was a long-standing latent bug, not a regression from the July 2026 `common.sh` platform-function unification or the `DEBUG\_TRACE` work; those changes only made the previously-silent abort **\*\*visible\*\*** (via the `EXEC:/EXIT:` breadcrumbs and xtrace).

**\*\*Fix (July 13, 2026):\*\*** Introduced `gateway\_should\_be\_running()` in `common.sh` and routed all four decision sites through it (`toggle.sh` main decision + `--restart` pre-check; `scripts/toggle-linux.sh` main decision + `--restart` pre-check). It decides from persisted intent:

1. `toggle.sh` captures the marker's existence into `GATEWAY\_MARKER\_AT\_STARTUP` **\*\*before\*\*** the top-of-script "defensive stale-marker clear" wipes it, and the helper honors that captured value (falling back to the live marker file for callers that run before the clear, e.g. the `--restart` pre-check).
2. Marker present STOP path; marker absent START path. This is instant and cannot hang or abort.
3. Only when the marker is absent **\*\*and\*\*** the Docker socket is actually reachable (`/var/run/docker.sock` or `~/colima/default/docker.sock` on macOS) does it consult `is\_gateway\_active` as a non-fatal reconciliation fallback (covers a marker lost to a `/tmp` wipe while containers are genuinely up). A dead socket short-circuits to "stopped" with no probe, so `set -e` can never be tripped.
4. The STOP path's `docker compose down` is now `|| true`-guarded so a disable always completes even with Docker down; the START path's `docker compose up` is deliberately left unguarded (a failed start **\*\*should\*\*** trigger cleanup). On Linux the marker isn't maintained, so the helper degrades to the `is\_gateway\_active`

fallback fast there because native Docker has no Colima VM to hang.

1932

1933 **\*\*How we got here (methodology note).\*\*** This one is worth recording because the *\*process\** mattered as much as the patch. The failure had been happening silently for a long time the log simply ended after a teardown with no error, so there was nothing to chase. We first attacked the *\*\*observability gap\*\** rather than guessing at the cause: we wrapped the heavyweight lifecycle commands (``colima start``, ``docker compose up/down``) so their output and exit codes always land in ``output.log``, added an ``EXIT: code= phase=`` breadcrumb that fires even when the process is killed, and put a full opt-in xtrace behind ``DEBUG_TRACE``. Only *\*then\** did we reproduce the failure and this time the trace said, unambiguously, ``EXIT: code=1 phase=start`` with the START body never entered. The instrumentation didn't fix anything, but it converted an invisible abort into a precise, reproducible fact. (Lesson, consistent with 15.12.18 and the LUKS post-mortem referenced in 15.11: *\*\*a mechanism that fails silently is worse than one that fails loudly\*\** so we make it fail loudly first.)

1934

1935 The root-cause leap, though, came from a *\*\*design-smell intuition, not the stack trace\*\**: the observation that it is nonsensical for a *\*disable\** to first check whether the gateway is *\*currently online\** a teardown should act on whether the gateway is *\*supposed\** to be running, not probe a subsystem to find out. That instinct turned out to name the bug exactly. The smell was a *\*\*conflated responsibility with an inverted dependency\*\**: to decide whether to *\*start\** the gateway, the code first had to *\*talk to\** the gateway (via Docker) the very thing START is responsible for bringing up so the check was guaranteed to be unavailable in precisely the cold-start case that mattered. A second tell was *\*\*false symmetry\*\**: `enable` (builds state) and `disable` (tears it down) are not symmetric operations, yet both ran the identical expensive probe; whenever two operations that should differ share suspiciously identical machinery, the code is usually lying about the structure of the problem, and reality eventually calls the bluff. The fix invented nothing the honest signal (``MARKER_FILE``, literally "the gateway is supposed to be up") already existed and was maintained in all the right places; it was simply never consulted for the decision it was made for. We just pointed the decision at the answer the codebase already knew. General principle worth keeping: *\*\*design smells here are load-bearing they tend to \*be\* latent bugs, not merely cosmetic\*\**, and are worth chasing on intuition even before a trace confirms them.

1936

1937 **\*\*Acceptance verified (static + live):\*\*** with marker present STOP (instant, even with Docker down); with marker absent + Docker down START (0 s, no ``set -e`` abort); a real run now proceeds past init to the gateway-status decision. ``bash -n`` clean on all three scripts; no new ``shellcheck`` findings; re-signed ``signatures`` (``toggle.sh``/``common.sh``/``toggle-linux.sh`` are in the signed set) and ``crypto.sh --verify`` exits 0.

1938

1939 **\*\*Follow-up (same day)** a ``set -e`` regression the fix introduced, and why only a live run caught it.**\*\*** The first cut of the marker capture wrote the intent as a full ``if [ -f "$MARKER_FILE" ]``; then `GATEWAY_MARKER_AT_STARTUP="yes";` else `GATEWAY_MARKER_AT_STARTUP="no";` fi. On macOS ``/bin/bash`` 3.2, with ``set -e`` active and the ``DEBUG_TRACE`` xtrace redirect (``exec 2> >(tee)``) in effect, that construct **\*\*** aborted the script at init**\*\*** the moment the marker was absent the false ``[ -f ]`` status leaked through the compound and ``set -e`` killed the run (``EXIT: code=1 phase=init``), before the gateway logic ran at all. Notably this did **\*\*not\*\*** reproduce in isolated ``bash -c`` snippets of the same block it only surfaced in the script's real runtime context so we had to bisect against the actual ``toggle.sh`` (swap the block for a trivial assignment the script sailed past init) to prove causation. Fix: write it as a default assignment plus an ``if`` **\*\*without an `else`\*\*** (`GATEWAY_MARKER_AT_STARTUP="no" then` if [ -f "$MARKER_FILE" ]``; then `GATEWAY_MARKER_AT_STARTUP="yes";` fi), which yields status 0 when the condition is false. Lesson reinforced: ``set -e`` on legacy bash is genuinely treacherous around conditionals, and the observability work paid for itself again the breadcrumb pinned the abort to init instantly, and a live run (not static checks) was the only thing that exposed a heisenbug invisible to isolated reproduction.

1940

1941 ##### 15.12.20 Honey-Port Tripwire Accumulation (July 14, 2026)

1942 **\*\*Symptom:\*\*** After many toggle-ons, ``ps`` showed a pile of idle ``bash ./toggle.sh`` subshells, each blocked on an ``nc -l localhost <port>`` listener one leaked per toggle-on, never reaped.

1943

1944 **\*\*Root Cause:\*\*** The Honey-Port Tripwire arms a fresh random port each START via a disowned ``( nc -l ) &`` subshell. ``nc -l`` blocks until a connection that (by design) almost never comes, so each toggle-on leaks one idle listener; because the port is random, successive listeners never collide and simply stack up over time.

1945

1946 **\*\*Fix:\*\*** ``toggle.sh`` now reaps the previous tripwire before arming a new one ``pkill -f "nc -l localhost" ( macOS ) / `pkill -f "nc -l .:127.0.0.1" + "nc -l localhost" (Linux) immediately before` HONEY_PORT=$( get_free_port )` at both arming sites. No `sudo` needed (the listeners are user-owned). Net effect: at most one tripwire alive at a time instead of one per toggle. (The same `pkill` pattern cleared the backlog that this session's ~15 test toggles produced.)`

1947

1948 ##### 15.12.21 Post-Wake Self-Feedback Loop: The Watcher Re-Fires On The Route Change It Caused (July 14, 2026)

1949 **\*\*Symptom:\*\*** After a container restart (or wake), ``output.log`` showed ``recover.sh --post-wake`` launching ~6 times in a row, one launch roughly every ~40s, each exiting ``exit=0`` recovery *\*worked\** every time yet kept re-firing "until it worked." Every launch was triggered by ``NET: changedKey State:/Network/Global/IPv4``.

1950

1951 **\*\*Root Cause** a self-feedback loop, NOT PF.**\*\*** ``scripts/watcher.sh`` Listener 2 watches ``State:/Network/Global/IPv4`` via ``scutil n.watch`` and launches ``recover.sh --post-wake`` on any change. But post-wake re-asserts the exit node (``recover.sh`` ~L367, ``tailscale set --exit-node``), and even though ``set`` is deliberately used instead of ``--reset`` to minimise churn (``recover.sh`` ~L360-366), **\*\*re-asserting the exit node re-writes the host default route and the default route IS` State:/Network/Global/IPv4`\*\*** So each run mutates the exact key the watcher triggers on. The self-triggered change lands ~15-45s after the run finishes; ``run_recover``()'s old 10s debounce (reset to the run's *\*finish\** time at the tail of the function) had long expired by then, so the watcher re-fired. The loop self-terminates only once ``tailscale set`` becomes a no-op (route already correct) no route change no re-trigger.

1952



```

1953 **Why the debounce couldn't stop it:** the loop period (~40s = one full post-wake run + settle) far exceeds the
      10s debounce, which therefore only ever caught macOS's rapid *duplicate* scutil notifications for a single
      change (the `~0.2s` doublets), not the self-trigger.
1954
1955 **Not PF:** post-wake never calls `pfctl` (the kill-switch is armed once, in `scripts/common.sh`, not per cycle
      ); no PF-enable line falls inside the loop window. The `No ALTQ support in kernel` warning is normal macOS
      pfctl noise (Apple ships pfctl without ALTQ compiled in), and `Could not capture PF reference token` is a
      separate benign ref-count bookkeeping wart both unrelated to the loop.
1956
1957 **Fix:** `scripts/watcher.sh` `run_recover()` now arms a **post-recovery settle cooldown** (`/tmp/nullexit-
      watcher.settle-until`, `NULLEXIT_POSTWAKE_SETTLE_SECONDS` default 45s) **after** each post-wake run, and
      skips launching while `now < settle-until`. Because it is armed only after a recovery runs, a genuine first
      wake/roam is still handled immediately; only the Global/IPv4 change the run caused itself is absorbed. A
      real new roam during the window is delayed at most `POSTWAKE_SETTLE_SECONDS`, an acceptable trade for
      killing the loop. Watcher-side only `recover.sh` is unchanged. Severity was low (self-correcting, healthy
      every iteration, no leak/outage; cost was wasted ~28s runs + log noise).
1958
1959 ### 15.12.22 FaceTime (and Real-Time P2P Media) Over the Double-Tunnel: Why It Breaks + the Split-Route Plan (
      July 15, 2026)
1960 **Symptom:** A FaceTime call placed while nullexit is up fails to connect.
1961
1962 **Diagnosis nullexit is NOT blocking it; every layer passes the traffic (all verified live):**
1963 - **DNS:** no Apple/FaceTime domain is sinkholed `dig` via the gateway (`100.100.21.8`) returns real Apple A/
      CNAME records, none `0.0.0.0`/`127.0.0.1`. The `` domains (`gateway.push.apple.com`, `fmn.apple.
      com`, `identityservices.apple.com`, `fdr.apple.com`) return empty on `1.1.1.1` too CNAME/non-A, not our
      block.
1964 - **routing-fix `FORWARD`:** the exit-path rule `-i tailscale0 -o tun0 -j ACCEPT` carries ~203K pkts / 73M
      bytes and accepts **everything**.
1965 - **Country / threat blocklists** (`block_il`, `block_kp`, `block_malicious`): **0** packets dropped.
1966 - **WARP (gluetun) egress:** `-A OUTPUT -o tun0 -j ACCEPT` + **A POSTROUTING **o tun0 -j MASQUERADE` all
      outbound UDP (any port) leaves and is SNAT'd through WARP.
1967
1968 **Root cause double symmetric NAT ending at Cloudflare WARP.** Path: `host Tailscale (MASQUERADE) WARP (
MASQUERADE) Cloudflare shared egress`. WARP is a CGNAT-style shared egress with **no inbound UDP mappings
**, and the double-MASQUERADE presents as **symmetric NAT** the exact condition that defeats FaceTime's
      STUN/ICE UDP hole-punching. No P2P path can form; the Apple-relay fallback is unreliable through the
      stacked tunnel, and the 12801200 MTU squeeze degrades media. This is inherent to routing real-time P2P
      media through WARP, **not** a firewall rule. (Zoom/WhatsApp **calls** hit the same wall.)
1969
1970 **Latent finding (separate bug, NOT yet fixed):** the `FORWARD` chain's **i tailscale0 -p udp --dport 443` (
QUIC) and **--dport 853` (DoT) `REJECT` rules are **shadowed** by the earlier **i tailscale0 -o tun0 -j
ACCEPT` (first-match wins) verified **0 pkts** on each. So the intended QUIC-block / force-through-AdGuard
(block DoT) on the exit path is **not enforced**: a mesh client can use DoT (853) to bypass AdGuard
filtering entirely. To actually enforce, those port-`REJECT` rules must be inserted **before** the
tailscale0 tun0` ACCEPT (or scoped **o tun0`). **TODO.**
1971
1972 **Fix plan (chosen: Option 2 narrow, always-on, opt-in split-route). DESIGNED / STAGED, not yet implemented in
code:**
1973 - **Mechanism:** add more-specific host routes for FaceTime's Apple `/16`s via the **physical** gateway
      (`172.17.0.1`/`en0`). Longest-prefix match beats Tailscale's exit-node `0.0.0.0/1` + `128.0.0.0/1` routes,
      so only those `/16`s leave **direct** while everything else stays double-tunneled. Mirrors `
      add_warp_bypass_routes`.
1974 - **Wiring:** `add_apple_split_routes()` / `remove_apple_split_routes()` in `common.sh`; called in toggle START
      (after exit-node assert) + `recover.sh --post-wake` (survives roam) + removed on STOP; guarded by an **opt
      -in `env` flag (default OFF)** so the privacy posture is unchanged unless enabled.
1975 - **Privacy tradeoff:** whatever `/16`s go direct see the **real ISP IP** (not WARP). Apple already ties
      traffic to the Apple ID, so Apple-direct is the accepted scope; user chose **narrow** (FaceTime-infra `/16`
      s only) over all-`17.0.0.0/8` to minimise the surface.
1976 - **CIDR list to be finalised from an empirical capture (PENDING, deferred at user request):** host-side `sudo
      tcpdump -n -i <exit-utun, e.g. utun6> net 17.0.0.0/8` during a real call, then extract distinct `/16`s.
      Known-likely from DNS today: `17.57/16` (APNs push/ringing), `17.248/16` (iCloud/FaceTime gateway),
      `17.253/16` (aapling push CDN), `17.157/16` (GSA auth). **The relay/STUN `/16`s the load-bearing media
      path are call-time-discovered and **MUST come from the capture;** routing signaling-only will leave media
      broken. Signed scripts are untouched until the capture yields the media ranges. See P2P note 16.4.
1977
1978 ---
1979
1980 # Part IV Observations, Changelog & TODO
1981
1982 ## 16. Unverified Observations
1983
1984 ### 16.1 AdGuard Returns Two Different Sinkhole IPs (Unverified)
1985
1986 > **Status: Observed but not formally verified. Needs a deeper audit of the rule-compiler sources to confirm.**
1987
1988 **Observation (July 10, 2026):** When querying AdGuard Home for blocked ad domains, some domains resolve to
      `0.0.0.0` and others resolve to `127.0.0.1`. Both are blocked, both cause connection failure, but the
      different responses were unexpected.
1989
1990 ---
1991 doubleclick.net 0.0.0.0 (blocked)

```



```

1992 ads.google.com 127.0.0.1 (blocked)
1993 pagead2.googlesyndication.com 0.0.0.0 (blocked)
1994 scorecardresearch.com 127.0.0.1 (blocked)
1995 ```
1996
1997 **Likely Explanation (Unverified):** There are three historical schools of thought on the "correct" DNS
    sinkhole response, and AdGuard faithfully preserves whichever the source blocklist used:
1998
1999 | Sinkhole IP | Convention | Origin |
2000 | ---|---|---|
2001 | `0.0.0.0` | AdBlock-style lists (`\|\domain`) | Modern DNS-level blocking; AdGuard's native format. Fails
    fast OS doesn't attempt a TCP connection. Works for both IPv4 and IPv6 without a separate rule. |
2002 | `127.0.0.1` | Hosts-file-style lists (`127.0.0.1 domain`) | 90s-era `/etc/hosts` tradition. Steven Black's
    hosts list (500k+ domains) still uses this. |
2003 | `NXDOMAIN` | Purist / Pi-hole default | Returns "domain doesn't exist." Semantically most accurate but
    requires browsers to handle NXDOMAIN gracefully. |
2004
2005 The dual-response behavior in this project is because `rule-compiler` pulls from both AdBlock-format sources (
    Hagezi, uBlock origin lists) and hosts-format sources (Spamhaus DROP, Steven Black, Feodo Tracker), and
    AdGuard returns whatever sinkhole IP the matching rule specifies.
2006
2007 **To Verify:** Run `docker compose exec tailscale cat /etc/hosts` to confirm no hosts-file rules are being
    injected at the container level, and check which specific AdGuard rule matched each domain via the AdGuard
    query log at `http://100.100.21.8:3000/#logs`.
2008
2009 ### 16.2 AGY CLI Appears to Generate Native macOS Notification Pop-ups (Unverified)
2010
2011 > **Status: Observed but source unverified. macOS security logs ruled out system-level origin.**
2012
2013 **Observation (July 10, 2026):** While running a DNS blocklist verification via the Antigravity CLI (`agy`), a
    command was proposed that included `malware.wicar.org` (a known-safe DNS blocklist test domain) in a `dig`
    loop. Shortly after, a macOS-style notification popup appeared describing a "malicious" or "blocked" event.
    The AGY CLI terminal session then closed.
2014
2015 **Investigation:**
2016 - **macOS XProtect / Gatekeeper logs:** Empty. Zero events in the 30-minute window around the incident.
2017 - **macOS Endpoint Security / System Policy logs:** Empty.
2018 - **Broad `log show` search for "malicious", "malware", "blocked":** Empty.
2019 - **Conclusion:** macOS itself did not generate the notification. No system security subsystem fired.
2020
2021 **Likely Explanation (Unverified):** The notification appears to have originated from **AGY CLI's own internal
    safety system**. AGY CLI is a native macOS application (not just a terminal process) and has access to the
    macOS Notification Center API (`UNUserNotificationCenter`). When its guardrails detected a domain
    containing the word "malware" (`malware.wicar.org`) in a proposed shell command, it likely:
2022 1. Blocked the command from executing
2023 2. Posted a native macOS-style notification via the Notification Center
2024
2025 This is surprising behavior for a CLI tool most terminal programs don't send GUI notifications but AGY CLI
    appears to bridge both worlds: it runs in a terminal but is packaged as a native app with full system API
    access. The notification would be indistinguishable from a system security alert at a glance.
2026
2027 **Notes:**
2028 - `malware.wicar.org` is the DNS/AV equivalent of the EICAR test file a deliberately benign domain that
    triggers blocklist/AV systems to verify they work. It was included in the test set to confirm the AdGuard
    blocklist was catching it. AGY's safety system is not aware of this distinction.
2029 - For future DNS blocklist testing, use only ad/tracking domains (e.g., `doubleclick.net`, `adnxs.com`) to
    avoid triggering AGY's guardrails.
2030
2031 ### 16.3 `system_profiler SPAirPortDataType` Longevity (Wi-Fi Detection)
2032 The Wi-Fi security detection used by `watcher.sh`'s `detect_lan_p2p_mode()` currently relies on `
    system_profiler SPAirPortDataType` (after the `airport` binary removal full history in 15.11.6). This
    works as of macOS Sequoia (15.x) without sudo and without redaction, but given Apple's ongoing Wi-Fi
    privacy clampdown, it may be locked down in a future release. If `security` comes back empty on a future
    macOS version while on Wi-Fi, detection silently falls back to `allow=false` (a safe default). The long-
    term fix would be a small Swift CoreWLAN helper compiled and bundled in the repo.
2033
2034 ### 16.4 Tailscale LAN P2P Succeeds Cross-AP But Fails Same-AP (Client Isolation + No NAT Hairpin)
2035
2036 **Observation (July 15, 2026):** On the residence Wi-Fi, a Galaxy S24 in the lobby reception 6 floors away,
    associated to a *different* AP establishes a **direct, low-ping** Tailscale P2P session to an iPhone 11.
    The *same* pair on the *same* AP / same floor / close range does **not** go direct. Counterintuitively, **
    farther apart = P2P works**. The effect shows for the iPhone 11 (a plain client) but **not** for S24macOS
    or S24the gateway container.
2037
2038 **Mechanism:**
2039 - **Same AP:** the AP enforces **client isolation** (L2 deviceblock on one radio). Both devices then
    share one public IP; STUN returns identical external mappings, so reaching each other would need the
    residence router to **NAT-hairpin** which most campus/consumer gear does not do direct P2P fails and
    falls back to a DERP relay.
2040 - **Different AP (floors away):** large residence networks **segment per AP/area into distinct VLANs/subnets**,
    so the two devices sit on different L3 networks routed through the campus core, where client isolation

```

does not apply and the STUN path works **\*\*without\*\*** hairpin clean direct P2P.

2041 - **\*\*Why macOS / the container are immune:\*\*** nullexit pins a **\*\*direct WireGuard endpoint\*\*** (Tailscale port  
`41642`, `direct 172.x`), so those peers already hold a negotiated direct path and do not depend on same-AP  
L2 discovery. A plain iPhone 11 has no pinned endpoint, so its path is entirely at the mercy of the Wi-Fi  
topology which is exactly why **\*it\*** exposes the isolation effect.

2042 **\*\*Relevance:\*\*** refines `watcher.sh`'s `detect\_lan\_p2p\_mode()` (which probes same-subnet AP isolation via a  
broadcast ping). The probe correctly flags same-AP isolation, but this observation adds that isolation is  
**\*\*per-AP\*\*** and cross-VLAN P2P can succeed even when same-AP fails so a single "LAN P2P allowed?" boolean  
under-describes a multi-AP network. No code change; documents expected behavior. Related: the FaceTime NAT  
analysis in 15.12.22 (same STUN/hairpin physics, different symptom).

2044 **### 16.5 Working Theory** iPhone Cross-Device Visibility Reaches Other Phones, Not macOS (Instagram Auto-Open,  
July 15, 2026)

2046 **\*\*Observation:\*\*** Opening Instagram on a Galaxy S24 (specifically via **\*\*Chrome\*\*** on the S24, not Samsung  
Internet) auto-opens a tab on a nearby MacBook, but **\*\*only** when the S24 is physically close\*\* to the Mac  
it does not happen when the phone is away. **\*\*Chrome** on the S24 has zero granted Android permissions\*\* (no  
Bluetooth, no location, no nearby-devices) this rules out Chrome itself performing any BLE/Nearby  
proximity scan, since Android enforces those permissions at the app level regardless of what the browser  
ships. This pushes the likely actor down to **\*\*Google Play Services\*\***, which holds system-level Bluetooth/  
nearby-device access independent of per-app grants and could be doing the proximity detection on Chrome's  
behalf (e.g. via Android's "Nearby" APIs) then handing the result to Chrome only for the actual tab-open.  
A purely local, short-range channel is still required to explain the proximity gating regardless of which  
component performs it; this also rules out any cloud-mediated mechanism (Chrome cloud tab sync, FCM push,  
Tailscale/WAN P2P per 16.4) since those are distance-independent. Live probes during the session (`  
system\_profiler SPBluetoothDataType` for a paired device, `dns-sd -B \_services.\_dns-sd.\_udp` for 8s) found  
**\*\*no** paired Bluetooth device and no Android/Google/Instagram mDNS service\*\* only the Mac's own Bonjour  
services (AirPlay, RFB, `\_companion-link`, Spotify Connect). Inconclusive: a BLE-advertisement-only channel  
(no formal pairing) would not show up in either probe.

2048 **\*\*User's working theory:\*\*** the iPhone (also on hand, on the same network/Apple ID as the Mac) makes **\*other**  
2049 devices connected to it\* even an Android phone visible to nearby devices, but this visibility channel  
reaches **\*\*other** phones more readily than it reaches macOS\*\*. I.e., the iPhone may be acting as a discovery  
bridge/relay between the S24 and the Mac rather than the S24 and Mac talking directly. Combined with the  
Play Services angle above, a refined version is: **\*\*Play Services** on the S24 detects proximity to an Apple  
mesh (iPhone and/or Mac both signed into iCloud/AWDL-visible) via some cross-ecosystem discovery Google has  
built\*\*, rather than the S24 and Mac negotiating directly.

2050 **\*\*Cross-ecosystem refinement:\*\*** this isn't Google-only. Apple runs an equivalent OS-level background BLE  
2051 scanning layer for its own Continuity stack (Handoff, Universal Clipboard, AirDrop discovery, Find My the  
latter using **\*every** nearby Apple device\*, not just the owner's, as a passive BLE relay), implemented as  
system daemons (`bluetoothd`/`sharingd`/`rapportd`, already implicated in the unrelated AWDL wedge bug at  
15.10) rather than anything gated by app-level permission. The key mechanism that makes the two ecosystems  
mutually visible: **\*\*BLE advertisement packets** are broadcast in the clear over open air\*\* Apple's  
Continuity BLE frames are semi-proprietary but unencrypted at the advertisement layer, so any BLE radio (  
including Android's, via Play Services once Android's own "Wi-Fi & Bluetooth scanning" system-level toggle  
is on) can passively observe "an Apple device is here" and react, without pairing or explicit permission  
from Apple's side. So both ecosystems maintain a BLE mesh-discovery layer that sits below their own per-app  
permission model, and each side's layer is passively readable by the other's radio which is a plausible  
bridge for a Play-Services-mediated Chrome trigger reacting to iPhone/Mac Continuity broadcasts.

2052 **\*\*Status:\*\*** Unverified. Not yet reproduced under an active capture the July 15 session ran the mDNS browse \*  
2053 before\* confirming the phone was actively broadcasting, so it's not a clean negative result. To confirm:  
capture `tcpdump -i en0 udp port 5353` (mDNS) **\*\*and** a raw BLE advertisement capture\*\* (e.g. `nRF Connect`  
or similar BLE scanner app on a third device, or `sudo btmon`-equivalent on Linux macOS does not expose  
raw BLE advertisement capture without a hardware sniffer) at the exact moment Instagram is opened on the  
S24 while close, with the iPhone present vs. absent, to isolate whether the iPhone's Continuity broadcasts  
are the trigger or the S24Mac channel exists independently of it. On the Android side, checking `Settings`  
Location Location Services Wi-Fi & Bluetooth scanning` and whether **\*\*Google Play Services\*\*** (not Chrome)  
holds Bluetooth/nearby-device/location permissions would confirm or kill the Play-Services-as-actor theory.  
Related: 16.4 (same physical-proximity family of P2P/discovery quirks, different mechanism and inverted  
distance relationship).

2054 **## 17. Changelog / Recent Updates**

2055 Reflects the current branch's state. Each entry is a one-line summary + commit hash; read the commit message (  
2056 and the linked Incident Log entry) for full detail. Full post-mortems live in 15.

2057 - **\*\*July 10, 2026** `fix(race): suppress WARP Watcher during post-wake force-recreate`\*\* Fixed a race where  
switching Wi-Fi caused the host IP to leak as the raw ISP IP (`132.x`) on every roam; `recover.sh --post-  
wake`'s warp force-recreate tripped the WARP Watcher into a nuclear shutdown. Fixed via `/tmp/nullexit-warp`  
`-inhibit.marker`. See 15.2.7.

2060 - **\*\*July 10, 2026** startup/roaming stabilization suite\*\* Pre-flight Tailscale race (15.5.3), STUN-block cold-  
boot false negative (15.5.4), restart DNS-build race via `wait\_for\_dhcp\_settle` (15.5.5), nuclear/post-wake  
concurrency lock (15.2.8), infinite auto-recovery loop marker leak (15.2.9), PF kill-switch bricking  
SOCKS5 fallback (15.7.3), `--reset` host logouts (15.7.4), missing-flags abort (15.7.5), discarded docker-  
exec errors (15.12.18), and the SNAT endpoint poisoning fix (15.7.1).

2061 - **\*\*July 6, 2026** `fix: resolve edge-case vulnerabilities and system state leaks`\*\* Security-review hardening  
suite:

2062 1. **\*\*DNS Watcher Interface Independence:\*\*** `start\_dns\_watcher` now detects the active interface via `get\_active\_service()` instead of hardcoding a `grep` for "Wi-Fi".

2063 2. **\*\*Robust Lock Verification:\*\*** `toggle.sh` lock-file logic uses `ps -p` verification to prevent permanent lockouts from a recycled PID.

2064 3. **\*\*Fail-Closed Crypto:\*\*** `crypto.sh` integrity check switched from `[ -x ]` to `[ -f ]` so it runs even if git strips the `+x` bit.

2065 4. **\*\*Strict iptables` Pinning:\*\*** `post-rules.txt` FORWARD-chain rules explicitly pin `tailscale0` tun0` both directions, preventing escape via `eth0` during WARP drops.

2066 5. **\*\*Docker Local Network Isolation:\*\*** `docker-compose.yml` binds exposed WARP ports (`1080`, `5354`) to `127.0.0.1` to prevent LAN access.

2067 6. **\*\*Routing Verification & Sudoers:\*\*** Added a `netstat -nr` default-route override check and ensured `/usr/bin/python3` is permitted in the sudoers template.

2068 - **\*\*July 4, 2026 `cc789c5`\*\*** Concurrency mutex on `toggle.sh` (`/tmp/nullexit-toggle.lock`); defensive `TS\_AUTHKEY` init under `set -u`; Go `go vet`/`go build` steps in CI; AdGuard log capping; macOS BSD `grep` fixes in `setup-common.sh`; `recover-linux.sh` quoting/PID fixes; removed redundant bypass-routes block in `recover.sh`.

2069 - **\*\*July 4, 2026 `fix: post-wake routing loop & destructive recovery`\*\*** (1) `add\_warp\_bypass\_routes` now accepts explicit interface args so post-wake stops routing WARP traffic back into `utun\*` (was an infinite loop killing internet on every wake). (2) Replaced `sudo route -n flush` with targeted deletions of `0.0.0.0/1`/`128.0.0.0/1` + WARP bypass routes, so recovery no longer destroys loopback/multicast routes and wedges `configd`/`mDNSResponder` until reboot.

2070 - **\*\*July 4, 2026 `feat: secure execution`\*\*** Restricted the blanket `/usr/bin/python3` NOPASSWD sudo to exactly `/usr/bin/python3 -I ../scripts/dns-proxy.py`; locked `dns-proxy.py` with `chown root:wheel` + `chmod 755`; added native auth prompts to `Toggle Gateway.app`/`Recover Gateway.app`; `toggle.sh` now captures `warp` logs on failure before teardown.

2071 - **\*\*July 4, 2026 Go ports + refactor\*\*** `logger` and `rule-compiler` ported to Go multi-stage Docker builds (15.8.6); ~200 lines of duplicated bash extracted to `scripts/common.sh` (15.9.4); `crypto.sh` HMAC-SHA256 script integrity added.

2072 - **\*\*July 3, 2026 `fix: resolve numerous system regressions`\*\*** Fixed truncated `/etc/sudoers.d/nullexit` line; temporary `tmp`+`os.replace()` atomic writes in `sync-rules.py` (later reverted, see below); `recover.sh` `ts\_args` quoting; ANSI `%b`/`-e` output in `diagnose-host-leak.sh`/`fix-docker-bridge-collision.sh`; `toggle-linux.sh` using `resolvectl dns` instead of macOS `networksetup`; AdGuard REST refresh POST targeting port `3000`; restored `START\_GATEWAY=true`; purged stale port `80` mapping; verified `output.log` in `.gitignore`.

2073 - **\*\*July 3, 2026 `unlock-files.sh` + revert atomic writes\*\*** Reverted all 5 `tmp`+`os.replace()` sites in `sync-rules.py` to direct writes; created `scripts/unlock-files.sh` to permanently fix stale `0444`/`000` file permissions via atomic rename (no `chmod`). Covers `.env`, `compiled\_rules.txt`, `ip\_blocklist.ipset`, `cache/\*.txt`, `cache/ip/\*.txt`, `data/filters/\*.txt`. See 15.9.6 + 3.

2074 - **\*\*July 3, 2026 Overnight leak + geo-blocking\*\*** Introduced `scripts/host-leak-probe.sh` (sub-second host-egress detector, 15.6.1) and the statistical-threshold WARP auto-shutdown watcher (15.6.2); fixed the overnight silent IP leak (15.6.3).

2075 - **\*\*July 2, 2026 `8c5f8e4e` + `181e1e1`\*\*** Two infinite routing loops fixed: host-VM tunnel loop (WARP bypass routes via physical gateway IP + `setup\_exit\_node\_routing` forcing default to `utun\*`) and Docker Compose subnet takeover (`10.200.1.0/24` lock). `--accept-routes=true` added explicitly to all `tailscale up --reset` calls. See 15.2.2.

2076 - **\*\*July 2, 2026 auto-recovery daemon\*\*** `scripts/watcher.sh` + launchd LaunchAgent documented. See 15.5.1.

2077 - **\*\*July 1, 2026 `c5b92c0` (refactor: personal-leak cleanup)\*\*** Eliminated every residual personal-username leak across source + config; launchd plist uses `WorkingDirectory` + relative program arg with the `\_\_NULLEXIT\_HOME\_\_` sentinel (install recipe gains a single `sed -i 's|\_\_NULLEXIT\_HOME\_\_|\$(pwd)|'` step); 8 devref prose sites genericized to `~/colima/...`, `~/Library/LaunchAgents/...`, `\$USER`, ``.

2078 - **\*\*July 1, 2026 `a5e2a5a` (refactor: script-relative paths)\*\*** Replaced 6 hardcoded `/Users/<user>/.../nullexit/...` literals with script-relative resolution. `recover.sh`/`watcher.sh` inject `SCRIPT\_DIR=\$(cd "\$(dirname "\${BASH\_SOURCE[0]}")" && pwd)`; `RECOVER=\$(SCRIPT\_DIR/..recover.sh)` resolves without launch-path dependency.

2079 - **\*\*June 28, 2026 `f8f8e4e` (M1: scripts/ move)\*\*** 8 internal scripts consolidated into `scripts/`; the 4 user-facing/orchestrator/config files (`toggle.sh`, `recover.sh`, `setup.sh`, `docker-compose.yml`) stayed at repo root.

2080 - **\*\*June 28, 2026 `0875ef3` (ci: glob)\*\*** CI `py\_compile` switched to `python3 -m py\_compile scripts/\*.py` after M1 left root with zero `.py` files.

2081 - **\*\*June 28, 2026 `25b37a1` (docs: scripts/ prefix)\*\*** `devref.md` + `README.md` updated to the `scripts/` prefix for all 8 moved files (~31 patches).

2082 **### End-to-end verification (commit `c5b92c0`)**

2083 Manual START STOP cycle ran clean on macOS after path relativization + personal-leak cleanup:

2084 - `bash toggle.sh` START: exit 0, ~135s (cold Colima boot); marker present (`2026-07-02T03:41:41Z`), all 5 containers healthy, host DNS hijacked to `100.100.21.8` (no fallback), exit node selected, Cloudflare trace through WARP succeeds.

2086 - `bash toggle.sh` STOP: exit 0, ~12s; marker cleared, all nullexit containers gone, host DNS restored to `1.1.1.1`, Tailscale disconnected.

2087 **## 18. TODO & Future Work**

2088

2089 **### 18.1 TODO**

2090

2091 \* **\*\*FaceTime / real-time P2P split-route (Option 2, narrow):\*\*** Implement `add\_apple\_split\_routes()` / `remove\_apple\_split\_routes()` in `common.sh` (route FaceTime's Apple `/16`'s direct via the physical gateway, opt-in `.env` flag, default OFF), wired into toggle START/STOP + `recover.sh --post-wake`. **\*\*Blocked on\*\*** an empirical `tcpdump` capture of the relay/STUN `/16`'s during a real call (user-triggered). Full design + diagnosis in 15.12.22.

2092 \* **\*\*Fix the shadowed exit-path port REJECTs (DoT/QUIC bypass):\*\*** The `FORWARD` `-i tailscale0 -p udp --dport 853`/`--dport 443` REJECT rules are shadowed by the earlier `tailscale0 tun0` ACCEPT (0 pkts), so the intended block-DoT / force-AdGuard and QUIC-block are **\*\*not enforced\*\*** a mesh client can bypass AdGuard

via DoT (853). Re-order the REJECTs before the ACCEPT (or scope `-o tun0`). See 15.12.22.

2093 \* **Decide toggle STOP/START from persisted \*intent\*, not a live probe:** `toggle.sh` decided STOP-vs-START by calling `is_gateway_active()`, which live-probes the running system (`docker compose ps` + DNS + SOCKS checks). When Colima/Docker was down the `docker compose ps` blocked for the full 15 s `run_with_timeout` window and, under `set -e` + the `ERR` trap, aborted the script *before* the START body ran so the path that would boot Colima never executed (`EXIT: code=1 phase=start`, seen 57 historically).~~ **Closed** added `gateway_should_be_running()` in `common.sh`, which reads persisted intent (`MARKER_FILE`) captured into `GATEWAY_MARKER_AT_STARTUP` before the defensive stale-marker clear, and only falls back to `is_gateway_active` when the marker is absent *and* the Docker socket is actually reachable (so a dead-socket probe can neither hang nor trip `set -e`). All four decision sites (`toggle.sh` main + `--restart`, `toggle-linux.sh` main + `--restart`) now use it, and the STOP `docker compose down` is `|| true`-guarded so a disable always completes even with Docker down. See 15.12 Misc Resolved.

2094 \* **Verify Wi-Fi Roaming:** Investigate and test whether switching Wi-Fi networks now successfully and smoothly recovers the gateway end-to-end without any tailscale flag errors or dead tunnels. (Pending user verification).

2095 \* **Fix Diagnostic Script Check 5/8:** Investigate why `diagnose-host-leak.sh` check 5/8 (`Host default route`) sometimes reports "default route goes via physical Wi-Fi Tailscale route assertion failed" on macOS even though the `warp-on` egress checks confirm that traffic is successfully tunneling.~~ **Closed** fixed by switching check 5 from `netstat -rn | grep default` to `route -n get 1.1.1.1`. macOS Tailscale uses interface-scoped routes and does not replace the DHCP default route entry. See 15.7.1 Known Unknowns.

2096 \* **Expose Tor Control Port (9051):** Consider mapping the Tor Control port to localhost and adding `nyx` (the Tor terminal status monitor) to allow users to inspect their entry, middle, and exit node IPs in real-time.~~ **Closed** mapped the Tor Control Port to a procedurally generated, localhost-bound port `TOR_CONTROL_PORT`, configured `ControlPort` and disabled cookie authentication in `docker/tor/entrypoint.sh`, and integrated it into the `/sweep` verification protocol.

2097 \* **Host-Only Tor Mode (Pluggable Egress):** Implement an `EGRESS_TYPE=tor_host_only` mode for users in heavily censored (DPI-restricted) environments. This mode would skip booting `warp` and `tailscale`, boot `adguardhome` and `tor` (using Telegram `obfs4` bridges), set AdGuard's upstream to Tor's `DNSPort`, and use the `pf` kill-switch to force all host Mac traffic strictly through the Tor network. This provides a 100% free, DPI-resistant evasion path without requiring an external VPS.

2098 \* **The Technical Realities (What You Need to Watch Out For):**

2099 \* **The UDP Problem:** Tor only transports TCP traffic. It does not support UDP (aside from DNS requests passed directly to its `DNSPort`). macOS is incredibly chatty over UDP (FaceTime, mDNS, QUIC protocols). Your `pf` rules defined in `scripts/pf.conf` must be configured to aggressively and silently DROP all host UDP traffic (except for local DNS queries to AdGuard). If you don't drop it cleanly, macOS services might hang indefinitely while waiting for timeouts.

2100 \* **The "Daily Driver" Friction:** Routing a whole host OS through Tor is brutal on performance. Background daemons (iCloud sync, macOS software updates, Spotlight indexing) will blindly attempt to download gigabytes of data over Tor, which could saturate the circuits or time out completely.

2101 \* **Bridge Rot:** State firewalls actively hunt and block obfs4 bridges. When a user's bridges burn out, they will lose all internet access due to the `pf` kill-switch. You will need to ensure the user has a frictionless way to update the `tor-bridges.txt` file and restart the Tor container without exposing their IP.

2102 \* **Exit Node Discrimination:** Because all traffic will exit from known Tor nodes, the user will face an onslaught of Cloudflare CAPTCHAs, and many banking or streaming services will flat-out reject the connections.

2103 **### 18.2 Future Work**

2104 - **Direct P2P on VM Hosts:** Non-Linux hosts run Docker inside a VM, making true UDP hole-punching *to* arbitrary external peers impossible; those peers fall back to the DERP relay unless bridged mode is enabled on a trusted LAN. The only fully general solution is a native Linux host (Raspberry Pi, Intel NUC) where Docker runs without VM hypervisor translation. *(The one case that works regardless is direct **hostcontainer** P2P the gateway and its own Mac, the same physical machine re-permitted via routing-fix `.host_ips` RETURN rules and confirmed `direct` even in plain user-mode networking; see 15.3.5 and 15.8.7.)*

2106 - **Native Linux Deployment:** Benchmark on a Raspberry Pi native container footprint is ~75MB without macOS hypervisor overhead.

2107 - **Mesh-Wide Filesystem (SFTP over Tailscale):** Any mesh device passively exposes its filesystem over SFTP. Android via Termux + OpenSSH + wakelock; iOS is client-only due to background process limits.

2108 - **Post-Quantum Cryptography (PQC):** Tailscale uses Curve25519, theoretically vulnerable to "harvest now, decrypt later" quantum attacks. A future iteration could replace the Tailscale container with raw WireGuard + [Rosenpass](https://rosenpass.eu/) to negotiate post-quantum PSKs.

2109 - **Egress Compartmentalization (Pluggable Architecture):** See 12.9 for the full plan to make the entire egress layer swappable via a single `.env` variable.

## 5 diagrams.md

```

1 # nullexit Flow Diagrams & Architecture
2
3 ---
4
5 ## 1. System Architecture What's Running & Where
6
7 ```mermaid
8 graph TB
9     subgraph INTERNET[" Internet / Cloudflare Edge"]

```

```

10     CF["Cloudflare WARP\n(Exit Point)\nwarp=on"]
11 end
12
13 subgraph TAILNET[" Tailscale Mesh (100.x.x.x)"]
14     PHONE[" Phone / Laptop\n(Tailnet Client)"]
15 end
16
17 subgraph MACHOST[" macOS Host"]
18     direction TB
19     PFFIREWALL["macOS pf Firewall\n(Kill-Switch & TCP MSS Clamping)"]
20     TSDAEMON["tailscaled\n(brew service)\nhost Tailscale"]
21     HOSTDNS["Host DNS\nGateway IP\n(hijacked by toggle.sh)"]
22     HOSTSOCKS["Host SOCKS5\n127.0.0.1:1080\n(fallback only)"]
23     HOSTTOR["Host Tor SOCKS5\n127.0.0.1:$TOR SOCKS_PORT\n(randomized)"]
24     HOSTTORCTRL["Host Tor Control\n127.0.0.1:$TOR_CONTROL_PORT\n(randomized)"]
25
26     subgraph MONITORS[" Background Daemons (toggle.sh children)"]
27         WARPWATCH["WARP Watcher\n(every 30s, via docker exec)\n output.log"]
28         DNSWATCHER["DNS Watcher\n(re-hijacks if drift detected)\n output.log"]
29         CAFFEINATE["caffeinate -i\n(sleep prevention)"]
30     end
31
32     subgraph COLIMAVM[" Colima VM (Linux, 600MB RAM)"]
33         subgraph NETNS["warp container network namespace (shared by ALL containers)"]
34             GLUETUN["warp / Gluetun\nWireGuard tun0\n(owns the namespace)"]
35             TS["tailscale container\nAdvertises exit node\ntailscale0 interface"]
36             SOCKS["tailscale\nSOCKS5\nport 1080"]
37             ADGUARD["adguardhome\nDNS sinkhole\nport 5335"]
38             ROUTINGFIX["routing-fix sidecar\nGeo-IP Blocking & Go Logger\n(writes blocked.log)"]
39             TOR["tor container\nSOCKS5: port 9050\nControl: port 9051\nTransparent Proxy: port 9040\n"]
40         end
41     end
42
43     subgraph LAUNCHD[" launchd (always-on)"]
44         WATCHER["scripts/watcher.sh\n(sleep/wake + Wi-Fi roam\n+ LAN P2P auto-detection)"]
45     end
46 end
47
48 subgraph RECOVERY[" Recovery"]
49     RECOVER["recover.sh\n(nuclear reset)"]
50     RECOVERPOST["recover.sh --post-wake\n(light refresh)"]
51 end
52
53 %% Traffic flow
54 PHONE -->|"WireGuard / Tailscale mesh"| TSDAEMON
55 TSDAEMON -->|"exit-node route utun*"| TS
56 TS -->|"FORWARD tun0 (MASQUERADE)"| GLUETUN
57 GLUETUN -->|"WireGuard UDP (macOS Host Egress)"| PFFIREWALL
58 PFFIREWALL -->|"TCP MSS Clamped / Kill-Switch check"| CF
59 TS -->|"198.18.0.0/15 PREROUTING redirect"| TOR
60 TOR -->|"internal path / exit nodes via tun0"| GLUETUN
61
62 %% DNS
63 HOSTDNS -->|"UDP:53 TCP:5354"| ADGUARD
64 ADGUARD -->|"DNS upstream via tun0"| CF
65 ADGUARD -->|"onion queries localhost:5353"| TOR
66
67 %% Monitoring
68 WARPWATCH -->|"docker exec warp wget cdn-cgi/trace"| GLUETUN
69 DNSWATCHER -->|"networksetup -getdnsservers"| HOSTDNS
70
71 %% Logging
72 ROUTINGFIX -->|"writes"| BLOCKEDLOG["blocked.log\n(Host-side)"]
73 HOSTTOR -->|"port mapping"| TOR
74 HOSTTORCTRL -->|"port mapping"| TOR
75
76 %% Recovery triggers
77 WARPWATCH -->|"6 consecutive warp=off"| RECOVER
78 WATCHER -->|"sleep/wake / roam"| RECOVERPOST
79 RECOVERPOST -->|"if gateway unhealthy"| RECOVER
80 WATCHER -->|"writes .lan_p2p_detected"| ROUTINGFIX
81
82 style MONITORS fill:#1a1a2e,color:#e0e0ff
83 style COLIMAVM fill:#0d2137,color:#e0e0ff
84 style NETNS fill:#0a1628,color:#e0e0ff
85 style RECOVERY fill:#2d0a0a,color:#ffe0e0
86 style INTERNET fill:#0a2d1a,color:#e0ffe0
87 style TAILNET fill:#1a2d0a,color:#e8ffe0
88 style LAUNCHD fill:#2d1a2d,color:#ffe0ff
89

```



```

90 ---
91 ---
92
93 ## 2. toggle.sh Full START Flow
94
95 ```mermaid
96 flowchart TD
97     START(["./toggle.sh"])
98     CHECK{"is_gateway_active?\n(containers running OR\nDNS 1.1.1.1)"}
99
100     START --> CHECK
101     CHECK -->|"YES already ON"| STOPFLOW[" See STOP Flow"]
102     CHECK -->|"NO start it"| S1
103
104     S1["1. Reset DNS 1.1.1.1\n(prevent deadlocks)"]
105     S2["2. tailscale down\n(prevent exit-node routing during boot)"]
106     S3["3. Check Colima VM\n(colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged)"]
107     S4["4. Configure VM swap\n(400MB, prevent OOM)"]
108     S5["5. Clean corrupted AdGuard config\n(prevent container crash loop)"]
109     S6["6. docker compose up -d --build\n(compile Go threat rules & boot containers)"]
110     S7["7. Poll tailscale status\n(wait for 'offers exit node'\nup to 60s)"]
111     TSREADY{"container\nTailscale ready?"}
112     ABORT([" ABORT\n cleanup_handler"])
113     S8["8. Resolve gateway Tailscale IP\n(.gateway_ip or docker exec)"]
114     S9["9. Verify tailscaled running\n(auto-start if needed)"]
115     S10["10. tailscale up --reset\n--exit-node=\n--accept-routes=true\n--ssh=true --accept-dns=false"]
116
117     PF1{"Pre-flight 1/3\ntailscale ping gateway"}
118     PF2{"Pre-flight 2/3\ndig +tcp AdGuard DNS"}
119     PF3{"Pre-flight 3/3\nWARP internet check"}
120     ALLPASS{"All 3 pass?"}
121
122     EXITPATH[" EXIT NODE PATH\n tailscale up --exit-node=\n force_dns_to_gateway (3 retry)\n SOCKS5 disabled"]
123     SOCKSPATH[" SOCKS5 FALLBACK PATH\n macOS SOCKS5 127.0.0.1:1080\n DNS proxy 127.0.0.1:53\n No exit node ( SKIP_EXIT_NODE=true)"]
124
125     DAEMONS["Start background daemons\n caffeinate -i (sleep prevention)\n DNS Watcher\n WARP Watcher\n(all write output.log)"]
126
127     RULECOMPILE["Rule compiler status\n(log rule/IP counts)"]
128     WRITEMARKER["write_gateway_active_marker\n(/tmp/nullexit-gateway-active.marker)"]
129     DONE([" Gateway UP"])
130
131     S1 --> S2 --> S3 --> S4 --> S5 --> S6 --> S7
132     S7 --> TSREADY
133     TSREADY -->|"NO (40 NoState)"| ABORT
134     TSREADY -->|"YES"| S8
135     S8 --> S9 --> S10
136     S10 --> PF1 --> PF2 --> PF3
137     PF1 & PF2 & PF3 --> ALLPASS
138     ALLPASS -->|"YES"| EXITPATH
139     ALLPASS -->|"NO"| SOCKSPATH
140     EXITPATH --> DAEMONS
141     SOCKSPATH --> DAEMONS
142     DAEMONS --> RULECOMPILE --> WRITEMARKER --> DONE
143
144     style ABORT fill:#5c1010,color:#fff
145     style EXITPATH fill:#0a3d1a,color:#c0ffc0
146     style SOCKSPATH fill:#3d2d00,color:#ffe0a0
147     style DAEMONS fill:#0a1f3d,color:#c0d8ff
148     style DONE fill:#0a3d1a,color:#fff
149 ---
150
151 ---
152
153 ## 3. toggle.sh STOP Flow
154
155 ```mermaid
156 flowchart TD
157     STOP(["./toggle.sh\n(gateway already ON)"])
158
159     ST1["1. tailscale down\n(disconnect from mesh)"]
160     ST2["2. docker compose down -t 5\n(stop all containers)"]
161     ST3["3. Reset DNS 1.1.1.1\n(networksetup on all services)"]
162     ST4["4. stop_local_dns_proxy\n(kill Python UDP proxy)"]
163     ST5["5. stop_pidfile_daemon\n(kill caffeinate, rm PID)"]
164     ST6["6. stop_pidfile_daemon\n(kill DNS Watcher, rm PID)"]
165     ST7["7. stop_warp_watcher\n(kill WARP Watcher, rm PID)"]
166     ST8["8. clear_gateway_active_marker\n(rm /tmp/nullexit-gateway-active.marker)"]
167     ST9["9. cleanup_network_state\n disable_all_proxies\n Flush DNS cache (dscacheutil)\n Flush stale routes\n Power-cycle Wi-Fi (offon)"]

```



```

168 DONE([" Gateway DOWN\nInternet restored"])
169
170 STOP --> ST1 --> ST2 --> ST3 --> ST4
171 ST4 --> ST5 --> ST6 --> ST7 --> ST8 --> ST9 --> DONE
172
173 style DONE fill:#0a3d1a,color:#fff
174
175 ---
176
177 ---
178
179 ## 4. Monitoring Layer The 3 Background Daemons
180
181 mermaid
182 graph LR
183     subgraph DAEMONS["Background Daemons (all start with gateway, stop on teardown)"]
184         direction TB
185
186         subgraph D1[" Sleep Prevention"]
187             CAFF["caffeinate -i\nPID: /tmp/nullexit-caffeinate.pid\n\nKeeps Mac awake while\ngateway is active"]
188         end
189
190         subgraph D2[" DNS Watcher"]
191             DNSW["Polls networksetup every ~30s\nPID: /tmp/nullexit-dns-watcher.pid\n\nIf DNS drifts from gateway IP\nre-hijacks automatically\noutput.log"]
192         end
193
194         subgraph D3[" WARP Watcher"]
195             WARPW["Polls every 30s via:\ndocker compose exec warp wget cdn-cgi/trace\nPID: /tmp/nullexit-warp-watcher.pid\n\nCounts consecutive warp=off\nWARP_FAIL_THRESHOLD (default 6=30s)\nruns recover.sh (nuclear)\noutput.log"]
196         end
197
198     end
199
200     subgraph BLIND["What each monitor can/can't see"]
201         B1["WARP Watcher sees:\nContainer WARP tunnel state\nHost NIC traffic\nFlash leaks < 5s"]
202         B2["Host-NIC blind spot covered by:\nPF kill-switch (kernel, fail-closed)\non-demand: diagnose-host-leak.sh\n(--watch to monitor) / sweep.sh"]
203     end
204
205     D3 -->|"covers"| B1
206     B1 -->|"gap closed by"| B2
207
208     style D1 fill:#1a2d0a
209     style D2 fill:#0a1a2d
210     style D3 fill:#2d0a2d
211     style BLIND fill:#1a1a1a,color:#aaa
212
213 ---
214
215 ---
216 ## 5. Traffic Flow Data Path (Normal Operation)
217
218 mermaid
219 sequenceDiagram
220     participant PHONE as Phone<br/>(Tailnet client)
221     participant TS_HOST as tailscaled<br/>(macOS host)
222     participant TS_CONTAINER as tailscale<br/>(container)
223     participant GLUETUN as Gluetun/warp<br/>(container)
224     participant CF as Cloudflare<br/>WARP Edge
225     participant INTERNET as Internet
226
227     Note over PHONE,INTERNET: Normal EXIT NODE path (all 3 pre-flights passed)
228
229     PHONE->>TS_HOST: WireGuard packet<br/>(encrypted, UDP 41641)
230     TS_HOST->>TS_CONTAINER: route via utun* tailscale0
231     TS_CONTAINER->>GLUETUN: FORWARD chain<br/>(iptables MASQUERADE on tun0)
232     GLUETUN->>TS_HOST: WireGuard tunnel (tun0) egresses macOS Host
233     Note over TS_HOST: macOS pf Firewall applies<br/>Kill-Switch & TCP MSS Clamping
234     TS_HOST->>CF: re-encrypted UDP packet (MSS 1160)
235     CF->>INTERNET: Plain HTTPS from<br/>Cloudflare IP
236
237     Note over PHONE,INTERNET: DNS Resolution path
238     PHONE->>TS_HOST: DNS query
239     TS_HOST->>TS_CONTAINER: UDP:53 TCP:5354 AdGuard :5335
240     TS_CONTAINER->>GLUETUN: AdGuard upstream DNS via tun0
241     GLUETUN->>CF: Encrypted DNS query
242     CF-->>GLUETUN: DNS response
243     GLUETUN-->>TS_CONTAINER: response

```

```

244 TS_CONTAINER-->>TS_HOST: response
245 TS_HOST-->>PHONE: DNS answer
246 ---
247 ---
248 ---
249
250 ## 6. recover.sh Decision Tree
251
252 mermaid
253 flowchart TD
254     INVOKE(["recover.sh called"])
255     MODE{"--post-wake\nflag?"}
256
257     subgraph POSTWAKE["--post-wake (light refresh, gateway stays UP)"]
258         PW1["Re-hijack DNS gateway IP"]
259         PW2["Flush DNS cache"]
260         PW3["Refresh tailscale exit-node"]
261         PW4["Force-recreate warp container\n(if Gluetun healthcheck failed)"]
262         PW5["Leave: containers, Wi-Fi,\ncaffeinate, DNS watcher untouched"]
263     end
264
265     subgraph NUCLEAR["Default (nuclear gateway torn down)"]
266         N0["Sudo keep-alive loop (background)"]
267         N1["1. tailscale down\n(disconnect from mesh)"]
268         N2["2. Reset DNS 1.1.1.1\n(ALL network services)"]
269         N3["3. disable_all_proxies\n(all interfaces)"]
270         N4["4. Flush DNS cache\n(dscacheutil -flushcache)"]
271         N5["5. Flush stale routes\n(WARP endpoint routes)"]
272         N6a["6a. docker compose down -t 30\n(stop all containers)"]
273         N6b["6b. stop_pidfile_daemon\n(kill caffeinate)"]
274         N6c["6c. stop_pidfile_daemon\n(kill DNS Watcher)"]
275         N7["7. Power-cycle Wi-Fi\n(networksetup offsleep 2on)"]
276         N8["8. Verify internet working\n(curl ifconfig.io)"]
277     end
278
279     DONE_PW(["Gateway refreshed\n(still running)"])
280     DONE_NUC(["Internet restored\n(gateway fully stopped)"])
281
282     INVOKE --> MODE
283     MODE -->|"yes"| PW1
284     PW1 --> PW2 --> PW3 --> PW4 --> PW5 --> DONE_PW
285     MODE -->|"no"| N0
286     N0 --> N1 --> N2 --> N3 --> N4 --> N5
287     N5 --> N6a --> N6b --> N6c --> N7 --> N8 --> DONE_NUC
288
289     style POSTWAKE fill:#0a2d1a,color:#c0ffc0
290     style NUCLEAR fill:#2d0505,color:#ffc0c0
291     style DONE_PW fill:#0a3d1a,color:#fff
292     style DONE_NUC fill:#3d1a00,color:#ffe0c0
293 ---
294 ---
295 ---
296
297 ## 7. Failure Paths & Self-Healing
298
299 mermaid
300 flowchart LR
301     subgraph EVENTS["Failure / Event"]
302         E1["WARP tunnel drops\n(warp=off)"]
303         E2["Mac wakes from sleep\nor Wi-Fi roams"]
304         E3["DNS drifts from\ngateway IP"]
305         E4["toggle.sh crashes /\nCtrl-C / SIGTERM"]
306         E5["Host-side leak\n(non-tunnel HOST NIC egress)"]
307         E6["Silent NAT timeout\n(ping stops working)"]
308     end
309
310     subgraph DETECTORS["Detector"]
311         D1["WARP Watcher\n(30s poll, docker exec)"]
312         D2["scripts/watcher.sh\n(launchd, always on)"]
313         D3["DNS Watcher\n(30s poll, networksetup)"]
314         D4["cleanup_handler\n(ERR/INT/TERM/HUP trap)"]
315         D5["PF kill-switch\n(kernel, always-on) +\nndiagnose-host-leak.sh --watch\n(on-demand)"]
316         D6["routing-fix.sh\n(30s curl over tun0)"]
317     end
318
319     subgraph ACTIONS["Response"]
320         A1["threshold consecutive off\nrecover.sh (nuclear)"]
321         A2["recover.sh --post-wake\n(gentle refresh)"]
322         A3["Re-hijack DNS\n(networksetup setdnsservers)"]
323         A4["Stop all daemons\nReset DNS 1.1.1.1\nTear down Tailscale"]
324         A5["Non-tunnel egress blocked\nfail-closed at kernel;\n--watch alerts on state change"]

```

```

325     A6["Blackhole internet traffic\n+ Drop TUNNEL_FAILED_CLOSED.marker"]
326     A7["Push macOS Desktop Notification\n(via watcher.sh listener)"]
327     end
328
329     E1 --> D1 --> A1
330     E2 --> D2 --> A2
331     E3 --> D3 --> A3
332     E4 --> D4 --> A4
333     E5 --> D5 --> A5
334     E6 --> D6 --> A6
335     A6 -->|"marker detected"| D2
336     D2 --> A7
337
338     style EVENTS fill:#2d1010,color:#ffcccc
339     style DETECTORS fill:#0a1a2d,color:#cce0ff
340     style ACTIONS fill:#0a2d0a,color:#ccffcc
341     ...
342
343     ---
344
345     ## 8. output.log Who Writes What
346
347     All events converge into a single `output.log` at the repo root.
348
349     | Writer | Event Types | Format |
350     |-----|-----|-----|
351     | `toggle.sh` | All startup/shutdown steps, errors, pre-flight results | Plain text with timestamps |
352     | `recover.sh` | Every recovery step (nuclear or post-wake) | Plain text |
353     | **WARP Watcher** | `WARP DOWN`, `WARP RECOVERED`, `WARP SHUTDOWN` | `[UTC timestamp]` prefix |
354     | **DNS Watcher** | DNS re-hijack events | Appended inline |
355     | `docker compose` | Container logs on failure (last 100 lines of warp) | Dumped on ERR path |
356     | `routing-fix.sh` | (via `nullexit-logger` inside container) | Structured |
357
358     **Grep cheatsheet:**
359     ```bash
360     grep 'WARP DOWN' output.log # Container-side tunnel drops
361     grep 'WARP SHUTDOWN' output.log # Auto-recovery triggered
362     grep 'WARP RECOVERED' output.log # Self-healed before threshold
363     grep -E 'EXEC:|EXIT ' output.log # Lifecycle breadcrumbs (last command + exit)
364     ```

```

## 6 toggle.sh

```

1  #!/bin/bash
2  # toggle.sh nullexit gateway toggle script (macOS + Linux)
3  # Automatically handles Docker containers, Colima, DNS hijacking, and Tailscale exit node routing.
4  #
5  # One file, both OSes: the dispatch below runs the self-contained Linux
6  # implementation (shuf / ss / nmcli / ip / systemd) and exits; on macOS it
7  # falls through to the macOS implementation (jot / netstat / networksetup,
8  # launchd / pf).
9
10 set -e
11
12 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
13
14 #
15 # LINUX TOGGLE self-contained; runs the full toggle then exits before the
16 # macOS body below. Native Linux tooling (shuf / ss / nmcli / ip / systemd).
17 # (Body inlined verbatim at column 0 it contains multiline `bash -c ""`
18 # strings that must NOT be re-indented.)
19 #
20 if [[ "$OSTYPE" != "darwin"* ]]; then
21 # Define log file (used by the lifecycle breadcrumb + run_logged helper)
22 LOG_FILE="$PWD/output.log"
23
24 # --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
25 # TOGGLE_PHASE is updated as the script moves through STOP/START so that if the
26 # process is killed (terminal closed, pkill, OOM) or exits on error, the EXIT
27 # trap records the final exit code AND the phase it died in making an
28 # interrupted START traceable in output.log instead of ending abruptly.
29 TOGGLE_PHASE="init"
30 _log_exit_breadcrumb() {
31     local rc=$?
32     echo "[$(date +%Y-%m-%d %H:%M:%S)] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
33     2>/dev/null || true

```

```

33 }
34 trap '_log_exit_breadcrumb' EXIT
35
36 # --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
37 # When enabled, mirror a full command-by-command xtrace to output.log. Off by
38 # default. Read directly from .env because common.sh (read_env_var) isn't
39 # sourced yet. Branch on bash version: BASH_XTRACEFD needs >= 4.1 (Linux
40 # usually has bash 5); fall back to teeing stderr on older bash. We never use
41 # the `exec {var}>>` dynamic-fd form (4.1+ only).
42 DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d '=' -f2- | tr -d '\n' | tr '[:upper:]' '[:lower:]' | tr -d '[:space:]')
43 if [ "$DEBUG_TRACE" = "true" ]; then
44     echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled xtrace mirrored to output.log (bash ${BASH_VERSIONINFO[0]}.${BASH_VERSIONINFO[1]})" >> "$LOG_FILE"
45     export PS4='+ [$(SECONDS)s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: ' # subshell-free clock: a $
46     () in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
47     if [ "${BASH_VERSIONINFO[0]}" -gt 4 ] || { [ "${BASH_VERSIONINFO[0]}" -eq 4 ] && [ "${BASH_VERSIONINFO[1]}" -ge 1 ];
48     }; then
49         exec 9>>"$LOG_FILE"
50         export BASH_XTRACEFD=9
51         set -x
52     else
53         # bash < 4.1: plain stderrfile redirect. NOT a process substitution under
54         # set -e, xtrace flowing through a backgrounded `tee` can abort the script
55         # (failed write/SIGPIPE trips set -e). See devref 15.11.8.
56         exec 2>>"$LOG_FILE"
57         set -x
58     fi
59 fi
60
61 # --- PROCEDURAL PORT GENERATION ---
62 # To prevent static port fingerprinting by malware, we randomize exposed ports
63 # on every boot and persist them to a hidden environment file.
64 PORTS_FILE="/tmp/nullexit-ports.env"
65 if [[ "$1" == "" || "$1" == "--restart" || "$1" == "restart" ]]; then
66
67     # Collision-avoidance function
68     GENERATED_PORTS=""
69     get_free_port() {
70         local port
71         while true; do
72             # Linux uses shuf or $RANDOM instead of jot
73             port=$(shuf -i 10000-65000 -n 1 2>/dev/null || echo $((10000 + RANDOM % 55000)))
74             # 1. Prevent self-collision within this loop
75             if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
76                 continue
77             fi
78             # 2. Prevent collision with ANY process on the machine (root daemons, terminal apps, etc)
79             # We use ss to read the kernel socket table directly.
80             if ! ss -ltn | awk '{print $4}' | grep -q ":$port"; then
81                 echo "$port"
82                 return
83             fi
84         done
85     }
86
87     export TOR SOCKS_PORT=$(get_free_port)
88     GENERATED_PORTS="$GENERATED_PORTS $TOR SOCKS_PORT"
89
90     export TOR_CONTROL_PORT=$(get_free_port)
91     GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"
92
93     export SOCKS_PROXY_PORT=$(get_free_port)
94     GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"
95
96     export DNS_PROXY_PORT=$(get_free_port)
97
98     echo "export TOR SOCKS_PORT=$TOR SOCKS_PORT" > "$PORTS_FILE"
99     echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
100     echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
101     echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
102 elif [ -f "$PORTS_FILE" ]; then
103     # On STOP, source the existing ports so docker-compose down matches
104     source "$PORTS_FILE"
105 else
106     # fallbacks if file is lost
107     export TOR SOCKS_PORT=9050
108     export TOR_CONTROL_PORT=9051
109     export SOCKS_PROXY_PORT=1080
110     export DNS_PROXY_PORT=5354
111 fi

```

```

110
111 # Start execution timer
112 TOGGLE_START_TIME=$SECONDS
113
114 if [[ "$1" == "restart" ]] || [[ "$1" == "--restart" ]]; then
115     echo "Executing Gateway Restart Sequence..."
116     # We must source common.sh here temporarily to check gateway state before we proceed
117     source "$SCRIPT_DIR/scripts/common.sh"
118     # Use persisted intent via gateway_state (echoes a string, never a return code;
119     # set -e-safe under set -x see devref 15.11.8). On Linux the marker is not
120     # maintained, so this degrades to the is_gateway_active fallback (fast native Docker).
121     if [ "$(gateway_state)" = "running" ]; then
122         echo "Gateway is currently running. Stopping it first..."
123         bash "$0"
124         echo "Gateway stopped. Now starting it..."
125     else
126         echo "Gateway is already stopped. Starting it..."
127     fi
128     bash "$0"
129     exit 0
130 fi
131
132 # Global flag to track if the script completed successfully
133 SUCCESS_RUN=false
134
135 # Global variable to track the currently running background command PID
136 CURRENT_BG_PID=""
137
138 # Source common bash functions
139 source "$SCRIPT_DIR/scripts/common.sh"
140
141
142 PID_FILE="/tmp/nullexit-cafeinate.pid"
143
144 # Start macOS caffeinate to prevent sleep while gateway is running
145 start_sleep_prevention() {
146     if ! command -v systemd-inhibit >/dev/null 2>&1; then
147         echo " [Warning] systemd-inhibit tool not found. Sleep prevention unavailable."
148         return 1
149     fi
150
151     # Stop any existing sleep prevention process first to avoid duplicates
152     stop_sleep_prevention
153
154     echo " Enabling system sleep prevention (systemd-inhibit)..."
155     # Run systemd-inhibit wrapped in a bash trap to prevent idle system sleep.
156     # Critically, this traps shutdown signals (SIGTERM) and automatically
157     # flushes hijacked DNS back to normal right before the PC powers off.
158     # We do not use recover.sh here because it is a heavy recovery script
159     # which takes too long and gets terminated before it can reset the DNS. Instead,
160     # we surgically and instantly restore normal DNS settings here.
161     nohup bash -c "
162         trap '
163             echo \"[Shutdown Trap] Reverting network settings...\\" >> \"\$PWD/output.log\" 2>&1
164             kill \$! 2>/dev/null || true
165             sudo -n resolvectl domain \"\$ACTIVE_SERVICE\" \"\" >> \"\$PWD/output.log\" 2>&1 || true
166             sudo -n resolvectl revert \"\$ACTIVE_SERVICE\" >> \"\$PWD/output.log\" 2>&1 || true
167             resolvectl domain \"\$ACTIVE_SERVICE\" \"\" >> \"\$PWD/output.log\" 2>&1 || true
168             resolvectl revert \"\$ACTIVE_SERVICE\" >> \"\$PWD/output.log\" 2>&1 || true
169             if command -v tailscale >/dev/null 2>&1; then
170                 tailscale up --accept-dns=false --exit-node= >> \"\$PWD/output.log\" 2>&1 &
171                 TS_DOWN_ARGS=\"\"
172                 if [ \"\$KILL_SWITCH\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
173                 tailscale down \$TS_DOWN_ARGS >> \"\$PWD/output.log\" 2>&1 &
174             fi
175             sleep 1
176             echo \"[Shutdown Trap] Cleanup complete.\\" >> \"\$PWD/output.log\" 2>&1
177             exit 0
178         ' SIGTERM SIGINT SIGHUP
179         systemd-inhibit --what=sleep --who=nullexit --why=\"gateway running\" --mode=block sleep infinity &
180         wait \$!
181         \" >> output.log 2>&1 &
182     local caffe_pid=$!
183     echo \"$caffe_pid\" > \"$PID_FILE"
184     echo " Sleep prevention active (PID $caffe_pid). Your PC won't sleep while the gateway is running."
185 }
186
187 # Stop sleep prevention when gateway is stopped
188 stop_sleep_prevention() {
189     if [ -f "$PID_FILE" ]; then
190         local caffe_pid

```

```

191     caffe_pid=$(cat "$PID_FILE")
192     if [ -n "$caffe_pid" ] && kill -0 "$caffe_pid" 2>/dev/null && ps -p "$caffe_pid" -o command= 2>/dev/null |
193         grep -q -E "systemd-inhibit|recover.sh"; then
194         echo " Stopping system sleep prevention (systemd-inhibit wrapper PID $caffe_pid)..."
195         kill "$caffe_pid" 2>/dev/null || true
196     fi
197     rm -f "$PID_FILE"
198 }
199
200 DNS_WATCHER_PID_FILE="/tmp/nullexit-dns-watcher.pid"
201
202 stop_dns_watcher() {
203     if [ -f "$DNS_WATCHER_PID_FILE" ]; then
204         local watcher_pid
205         watcher_pid=$(cat "$DNS_WATCHER_PID_FILE")
206         if [ -n "$watcher_pid" ] && kill -0 "$watcher_pid" 2>/dev/null; then
207             echo " Stopping background DNS Watcher (PID $watcher_pid)..."
208             kill -9 "$watcher_pid" 2>/dev/null || true
209         fi
210         rm -f "$DNS_WATCHER_PID_FILE"
211     fi
212 }
213
214 # Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
215 cleanup_handler() {
216     local exit_code=$?
217     local trigger_type="$1"
218     local line_no="$2"
219
220     if [ "$SUCCESS_RUN" = "false" ]; then
221         echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1..."
222         if [ -n "$ACTIVE_SERVICE" ]; then
223             reset_dns
224         fi
225
226         if [ -n "$CURRENT_BG_PID" ]; then
227             echo "Terminating active command (PID $CURRENT_BG_PID)..."
228             kill -15 "$CURRENT_BG_PID" 2>> output.log || true
229         fi
230
231         # Disconnect host Tailscale and close the GUI application on failure
232         disconnect_tailscale_host
233
234         # Stop local DNS proxy if running
235         stop_local_dns_proxy
236
237         # Stop sleep prevention
238         stop_sleep_prevention
239         stop_dns_watcher
240
241         # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
242         if [ -n "$ACTIVE_SERVICE" ]; then
243             cleanup_network_state
244         fi
245
246         if [ "$trigger_type" = "ERR" ]; then
247             echo -e "\n=====
248             echo "ERROR: Script failed at line $line_no with exit code $exit_code."
249             echo "=====
250             echo "Troubleshooting steps:"
251             echo "1. Run 'colima status' to verify the VM is active."
252             echo "2. Run 'docker compose ps' to check container health."
253             echo "3. Run 'docker compose logs' to view application logs."
254             echo "4. Run 'tailscale status' to check host Tailscale status."
255             echo "=====
256         fi
257     fi
258 }
259
260
261 trap 'cleanup_handler ERR $LINENO' ERR
262 trap 'cleanup_handler INT' INT
263 trap 'cleanup_handler TERM' TERM
264 trap 'cleanup_handler HUP' HUP
265
266 # NOPASSWD sudo
267 # The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
268 # needed (resolvectl, ip, systemctl, systemd-inhibit, pkill, and
269 # python3 scripts/dns-proxy.py for the fallback DNS proxy).

```



```

270 # All privileged calls below use `sudo -n` which will run without prompting.
271 # No credential caching loop is needed.
272
273 # Add Homebrew and standard paths
274 export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
275
276 # Check for local DNS proxy tools (used when tailnet data plane is unavailable)
277 # Uses a Python DNS proxy (handles TCP wire format correctly).
278 # Python3 is built into macOS no external dependencies.
279 DNS_PROXY_BIN=""
280 if command -v python3 >/dev/null 2>&1; then
281     DNS_PROXY_BIN="python3 -I"
282 elif command -v python >/dev/null 2>&1; then
283     DNS_PROXY_BIN="python -I"
284 fi
285
286 # Path to the DNS proxy script (sibling of this script)
287 DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/dns-proxy.py"
288
289 # Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
290 # (`brew install tailscale` + `systemctl start tailscaled`) NOT the
291 # Tailscale.app GUI so there is no .app bundle or Network Extension sandbox
292 # to launch, no menu-bar click to perform, and no scutil Network Service to
293 # start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
294 # if you ran the install with sudo) and the control socket is always available.
295 TS_BIN=""
296 if command -v tailscale >> output.log 2>&1; then
297     TS_BIN="tailscale"
298 fi
299
300 # Baseline: disable Tailscale DNS management at the daemon level.
301 # `tailscale set` writes a persistent preference. However, `tailscale up --reset`
302 # (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
303 # re-enable `--accept-dns=true` and nullify this `set`.
304 # Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
305 # `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
306 #
307 # This initial `set` is still useful as a baseline for non-reset operations
308 # (e.g. if the user runs `tailscale up` manually without --reset) and covers
309 # the brief window between this line and the first `--reset` call.
310 #
311 # Runs once at script start while tailscaled is still connected (if it was
312 # left running from a previous toggle), so the pref takes effect before any
313 # disconnection or reconnection logic.
314 if [ -n "$TS_BIN" ]; then
315     $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
316 fi
317
318 # Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
319 # *entire* teardown no need to mess with system network managers.
320 # tailscaled keeps running (as a systemd service), which is intentional:
321 # the next `tailscale up` then reconnects in seconds rather than waiting for a
322 # slow daemon launch. The `tailscale down` call also retains the user's
323 # auth state, so we don't force a browser re-login every cycle.
324 #
325 # Defined early (before the if/else branches and referenced by cleanup_handler's
326 # ERR trap) so that any failure before the main logic can still tear down Tailscale.
327 restart_tailscaled_daemon() {
328     echo " [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
329     echo " [Tailscale Recovery] Attempting to restart tailscaled service via systemctl..."
330     if run_with_timeout 15 sudo -n systemctl restart tailscaled >> output.log 2>&1; then
331         echo " [Tailscale Recovery] Successfully restarted tailscaled service."
332     else
333         echo " [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
334     fi
335     sleep 3
336 }
337 disconnect_tailscale_host() {
338     if [ -n "$TS_BIN" ]; then
339         echo "Disconnecting host Tailscale from mesh (tailscaled stays running as system service)..."
340         local ts_args=""
341         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
342         if ! run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1; then
343             restart_tailscaled_daemon
344             run_with_timeout 10 $TS_BIN down $ts_args >> output.log 2>&1 || true
345         fi
346     fi
347 }
348
349 ACTIVE_SERVICE=$(get_active_service)
350

```

```

351 # Helper to nuke leftover network state after teardown (proxy settings, routing
352 # table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
353 # problem where stale state kills internet after the gateway stops.
354 cleanup_network_state() {
355     echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."
356
357     # 1. Disable any leftover proxy settings (needs root on many macOS versions)
358     # (Skipped for Linux since we don't set global SOCKS proxy)
359     echo " Proxies disabled."
360
361     # 2. Flush DNS cache
362     if command -v resolvectl >> output.log 2>&1; then
363         sudo -n resolvectl flush-caches 2>> output.log || true
364     elif command -v systemd-resolve >> output.log 2>&1; then
365         sudo -n systemd-resolve --flush-caches 2>> output.log || true
366     fi
367     echo " DNS cache flushed."
368
369     # 3. Flush stale routing table entries
370     if command -v ip >> output.log 2>&1; then
371         sudo -n ip route flush cache >> output.log 2>&1 || true
372     fi
373     echo " Routing table flushed."
374
375     # 4. Power-cycle Wi-Fi to clear any lingering interface state
376     echo " Restarting network interface..."
377     sudo -n ip link set dev "$ACTIVE_SERVICE" down 2>> output.log || true
378     sleep 1
379     sudo -n ip link set dev "$ACTIVE_SERVICE" up 2>> output.log || true
380     echo " Interface power-cycled ($ACTIVE_SERVICE)."
381     sleep 3
382
383     echo "Network state cleanup complete."
384 }
385
386 # Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
387 # during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
388 # `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
389 # would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
390 # Capture current DNS before resetting so we can check if it was hijacked
391 INITIAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head
392 -1 || true)
393
394 echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
395 reset_dns
396
397 # Local Network Surveillance Check
398 echo -e "\n"
399 echo "Local Network Surveillance Check"
400 echo ""
401 # Count populated ARP cache entries to detect network scanning or high congestion
402 if command -v ip >> /dev/null 2>&1; then
403     ARP_COUNT=$(ip neigh | grep -iv 'incomplete' | wc -l | tr -d ' ')
404 else
405     # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
406     ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
407 fi
408 if [ "$ARP_COUNT" -gt 15 ]; then
409     echo -e " \033[1;33m[!] WARNING: High ARP activity detected ($ARP_COUNT devices in cache).\033[0m"
410     echo -e " \033[1;33m[!] This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)
411     .\033[0m"
412     echo -e " [] Your traffic remains fully encrypted and invisible.\n"
413 else
414     echo -e " [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
415 fi
416
417 echo "Checking Gateway Status..."
418 # Decide STOP vs START from persisted intent via gateway_state, which echoes a
419 # string (never a return code) to stay set -e-safe under set -x. See devref 15.11.8.
420 # On Linux this falls back to is_gateway_active (fast native Docker).
421 if [ "$(gateway_state)" = "running" ]; then
422     TOGGLE_PHASE="stop"
423     echo -e "\n=====
424     echo "Gateway is RUNNING. Stopping it now..."
425     echo -e "=====\\n"
426
427 # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
428 disconnect_tailscale_host
429 echo ""

```

```

430 echo "Stopping Docker containers..."
431 # `|| true`: a disable must always complete even if the Docker daemon is down.
432 run_logged docker compose down -t 5 || true
433
434 echo -e "\nDocker daemon handles container teardown automatically."
435
436 # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
437 # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
438 # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
439 echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
440 reset_dns
441
442 # 3. Stop local DNS proxy if running
443 stop_local_dns_proxy
444
445 # Stop sleep prevention
446 stop_sleep_prevention
447 stop_dns_watcher
448
449 # 4. Nuke leftover network state so internet actually works after teardown
450 cleanup_network_state
451
452 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
453 echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
454 else
455 TOGGLE_PHASE="start"
456 echo -e "\n=====
457 echo "Gateway is STOPPED. Starting it now..."
458 echo -e "=====
459
460 # 1. Prevent host exit-node deadlock during VM / Container startup
461 disconnect_tailscale_host
462
463 echo -e "\nUsing native Linux Docker..."
464
465 # 4. Clean up corrupted AdGuardHome configurations
466 if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
467     rm -f "adguard/conf/AdGuardHome.yaml"
468     echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
469 fi
470
471
472 # 5. Start compose services
473 write_host_ips
474 # Procedurally generated ports have already been exported at the top of the script.
475 echo " [Security] Procedurally randomized proxy ports generated to evade fingerprinting."
476
477 # --- HONEY-PORT TRIPWIRE ---
478 # Generate a random trap port. If any malware does a sequential port scan on localhost,
479 # it will inevitably hit this port. When it does, we instantly fire a desktop alert.
480 # Reap the honey-port from any previous toggle before arming a new one, so the
481 # `nc -l` tripwire never accumulates one idle listener per toggle-on (15.12.20).
482 # `nc -l` tripwire never accumulates one idle listener per toggle-on (15.12.20).
483 pkill -f "nc -l .*127.0.0.1" 2>/dev/null || true
484 pkill -f "nc -l localhost" 2>/dev/null || true
485 HONEY_PORT=$(get_free_port)
486 (
487     if nc -l -p "$HONEY_PORT" 127.0.0.1 >/dev/null 2>&1 || nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
488         echo "[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application
489             just pinged the local Honey-Port ($HONEY_PORT)!" >> "$PWD/output.log"
490         if command -v notify-send >/dev/null 2>&1; then
491             notify-send -u critical "nullexit OPSEC Alert" "An unknown application is aggressively scanning your
492             local ports. Malware port-scan suspected."
493         fi
494     fi
495 ) &
496 disown %1 2>/dev/null || true
497 echo " [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."
498
499 echo -e "\nStarting Docker containers..."
500 run_logged docker compose up -d
501
502 # Clean up the one-off rule compiler container so it doesn't clutter the Docker UI
503 docker compose rm -s -f rule-compiler >> output.log 2>&1
504
505 # 6. Wait for the gateway container's Tailscale connection to be ready
506 BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
507 USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')
508
509 echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
510 echo " (This can take 30-60s Tailscale goes through: NoState Starting Running Online)"

```

```

509 connected=false
510 consecutive_failures=0
511 for i in {1..60}; do
512     # Run tailscale status and capture its output so the user can see what's happening.
513     # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
514     ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)
515
516     if echo "$ts_output" | grep -q "offers exit node"; then
517         connected=true
518         break
519     fi
520
521     # Track consecutive failures to detect a stuck state.
522     # If we see 40+ consecutive 'NoState' responses, give up early
523     # the daemon is alive but can't connect to the coordination server.
524     if echo "$ts_output" | grep -q "NoState"; then
525         consecutive_failures=$((consecutive_failures + 1))
526         if [ "$consecutive_failures" -ge 40 ]; then
527             echo ""
528             echo " [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
529             echo " Tailscale daemon is running but can't reach the coordination server."
530             break
531         fi
532     else
533         consecutive_failures=0
534     fi
535
536     # Print a condensed status every 10 seconds so the user can see progress
537     if [ $((i % 10)) -eq 0 ]; then
538         ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
539         echo " [attempt $i/60] $ts_short"
540     elif [ $((i % 5)) -eq 0 ]; then
541         echo -n "$[i]"
542     else
543         echo -n "."
544     fi
545     sleep 1
546 done
547 echo ""
548
549 if [ "$connected" = "true" ]; then
550     echo "Gateway container is online on the Tailscale mesh."
551 elif [ "$BYPASS_PING" = "true" ]; then
552     echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
553 else
554     echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
555     echo " The container was seen in the following states during the wait:" >&2
556     echo "$ts_output" | sed 's/~/ /' >&2
557     echo "" >&2
558     echo " This could mean:" >&2
559     echo " 1. Tailscale auth key has expired check TS_AUTHKEY in .env" >&2
560     echo " 2. Network issue inside the container check: docker compose logs tailscale" >&2
561     echo " 3. gluetun VPN tunnel not connecting check: docker compose logs warp | tail -20" >&2
562     echo "" >&2
563     echo " Diagnose manually:" >&2
564     echo " docker compose exec -T tailscale tailscale status" >&2
565     echo " docker compose exec -T tailscale tailscale netcheck" >&2
566     cleanup_handler ERR $LINENO
567     exit 1
568 fi
569
570 # 6b. Resolve the gateway's Tailscale IP.
571 # Dynamically query the container for the IP (handles IP changes after re-auth)
572 # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
573 #
574 # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
575 # returns (\r) into captured output. If $TS_IP ends with \r, then
576 # `networksetup -setdnsservers` silently rejects the malformed IP and
577 # DNS stays at 1.1.1.1 the primary bug this project was hitting.
578 # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
579 # the first field of the first line.
580 TS_IP=""
581 if [ "$connected" = "true" ]; then
582     TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
583         awk 'NR==1{print $1; exit}' || true)
584     if [ -n "$TS_IP" ]; then
585         echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
586         echo "$TS_IP" > .gateway_ip
587     else
588         echo " [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
589         TS_IP=$(read_adguard_ip || true)

```

```

589 fi
590 else
591     TS_IP=$(read_adguard_ip || true)
592 fi
593
594 # 7. Connect host Mac to Tailscale mesh.
595 # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
596 # system LaunchDaemon via 'systemctl start tailscaled'), the previous five-
597 # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
598 # no menu bar item to click, and no Accessibility permission required.
599 #
600 # CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
601 # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
602 # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
603 # are impossible without the host on the mesh).
604 #
605 # We connect in two phases:
606 # Phase A Join mesh WITHOUT exit node (so pre-flight checks can run)
607 # Phase B Set exit node only after pre-flight checks confirm the gateway works
608 HOST_ON_MESH=false
609 SKIP_EXIT_NODE=false
610 if [ -n "$TS_BIN" ]; then
611     echo -n "Verifying tailscaled is reachable"
612     daemon_ready=false
613     for i in {1..10}; do
614         # `tailscale status` returns exit code 1 when the daemon is alive but
615         # disconnected ("Tailscale is stopped."). We check the daemon socket
616         # directly as a fallback if the socket exists, tailscaled is running
617         # even if the host isn't connected to the mesh yet.
618         if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
619             daemon_ready=true
620             break
621         fi
622         echo -n "."
623         sleep 1
624     done
625     echo ""
626
627     if [ "$daemon_ready" != "true" ]; then
628         echo -e "\033[0;33m[!] tailscaled daemon is not running. Attempting to start it...\033[0m"
629
630         # Try system-wide daemon (with timeout sudo may prompt for password)
631         if [ "$daemon_ready" != "true" ]; then
632             if run_with_timeout 10 sudo systemctl start tailscaled 2>> output.log; then
633                 echo "Started tailscaled as system LaunchDaemon."
634                 for i in {1..15}; do
635                     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
636                         daemon_ready=true
637                         break
638                     fi
639                     sleep 1
640                 done
641             fi
642         fi
643     fi
644
645     if [ "$daemon_ready" = "true" ]; then
646         echo "tailscaled is responsive (running as a system service)."
647
648         # Phase A: Join mesh without exit node
649         echo "Connecting host to Tailscale mesh (no exit node yet)..."
650         if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node=; then
651             HOST_ON_MESH=true
652             echo "Host is on Tailscale mesh."
653         else
654             echo -e "\033[0;33m[!] tailscale up failed even though the daemon is running.\033[0m"
655             echo "This usually means the host hasn't authenticated with your tailnet yet."
656             echo "Run this command in a terminal to authenticate:"
657             echo ""
658             echo "sudo tailscale up"
659             echo ""
660             echo "A browser will open. Log in to your Tailscale account, then re-run toggle.sh."
661         fi
662     else
663         echo -e "\033[0;31m[!] tailscaled could NOT be started automatically.\033[0m"
664         echo "Run this in a terminal to start and authenticate Tailscale:"
665         echo ""
666         echo "sudo systemctl start tailscaled"
667         echo "sudo tailscale up"
668         echo ""
669         echo "Then re-run toggle.sh."

```

```

670 fi
671 else
672     echo -e "\n[Warning] tailscale CLI not found on PATH. Skipping host Tailscale configuration..."
673 fi
674
675 # Phase B: Pre-flight checks + enable exit node only if safe
676 if [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" != "false" ] && [ -n "$TS_IP" ]; then
677     echo -e "\n"
678     echo "Pre-flight connectivity check"
679     echo ""
680
681     # Check 1: Can we reach the gateway via Tailscale (uses tailscale ping, not ICMP)?
682     # NOTE: tailscale ping exits 1 when the connection goes through a DERP relay
683     # ("direct connection not established") even though the pong was received.
684     # We check for 'pong' in the output (stderr discarded) to accept relayed connections.
685     echo -n " [1/3] Gateway reachable via Tailscale... "
686     ping_ok=false
687     # The host tailscaled needs a few seconds to propagate the new network map and
688     # establish a connection route after 'tailscale up --reset'. We retry up to 45 times.
689     for attempt in {1..45}; do
690         if $TS_BIN ping --until-direct=false -c 2 --timeout 2s "$TS_IP" 2>/dev/null | grep -q "pong"; then
691             ping_ok=true
692             break
693         fi
694         sleep 1
695     done
696     if [ "$ping_ok" = "true" ]; then
697         echo "PASS"
698     else
699         echo "FAIL"
700         SKIP_EXIT_NODE=true
701     fi
702
703     # Check 2: Can we reach AdGuard via localhost (exposed port ${DNS_PROXY_PORT})?
704     echo -n " [2/3] AdGuard DNS via localhost... "
705     if command -v dig >/dev/null 2>&1; then
706         if dig +tcp @127.0.0.1 -p ${DNS_PROXY_PORT} google.com +short +timeout=5 >/dev/null 2>&1; then
707             echo "PASS"
708         else
709             echo "FAIL"
710             SKIP_EXIT_NODE=true
711         fi
712     elif command -v nslookup >/dev/null 2>&1; then
713         if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 >/dev/null 2>&1; then
714             echo "PASS"
715         else
716             echo "FAIL"
717             SKIP_EXIT_NODE=true
718         fi
719     else
720         echo "SKIP (dig/nslookup not available)"
721     fi
722
723     # Check 3: Does the WARP container have external internet?
724     echo -n " [3/3] WARP container internet... "
725     if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace &>/dev/null;
726     then
727         echo "PASS"
728     else
729         echo "FAIL"
730         SKIP_EXIT_NODE=true
731     fi
732
733     # Act based on check results
734     if [ "$SKIP_EXIT_NODE" != "true" ]; then
735         echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
736         if $TS_BIN up --reset --ssh=true --accept-dns=false --exit-node="$TS_IP" --exit-node-allow-lan-access=
737         true; then
738             echo "Exit node enabled."
739         else
740             echo "[Warning] Failed to set exit node."
741             SKIP_EXIT_NODE=true
742         fi
743     fi
744
745     if [ "$SKIP_EXIT_NODE" = "true" ]; then
746         echo -e "\n033[0;33m[!] Pre-flight checks failed EXIT NODE + DNS HIJACK SKIPPED.\033[0m"
747         echo " Host stays on Tailscale mesh but your internet routes through your normal connection."
748         echo " Troubleshoot with:"
749         echo "     docker compose logs warp | tail -20"
750         echo "     docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"

```



```

749     echo ""
750     echo " Falling back to SOCKS5 proxy for traffic routing..."
751 fi
752 elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
753     echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
754     SKIP_EXIT_NODE=true
755 fi
756
757 # 8. Apply DNS Hijacking.
758 # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
759 # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
760 # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
761 if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
762     echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
763     echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
764     if force_dns_to_gateway "$TS_IP"; then
765         echo "DNS hijack verified: single server = $TS_IP."
766         start_dns_watcher "$TS_IP"
767     else
768         echo "[Warning] DNS hijack could not be verified."
769         echo " If DNS is broken, run manually:"
770         echo "     networksetup -setdnsservers \"$ACTIVE_SERVICE\" $TS_IP"
771         echo "     networksetup -setsearchdomains \"$ACTIVE_SERVICE\" ts.net"
772     fi
773 elif [ -n "$TS_IP" ]; then
774     echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
775     echo " Trying local DNS proxy for ad-blocking..."
776     if start_local_dns_proxy; then
777         echo " Ad-blocking active via local DNS proxy through AdGuard."
778     else
779         echo " Local DNS proxy unavailable. DNS stays at 1.1.1.1."
780         echo " AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
781     fi
782
783 # Enable SOCKS5 proxy for traffic routing through WARP
784 # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
785 echo -e "\n Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
786 echo " (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
787 enable_socks_proxy "$ACTIVE_SERVICE"
788
789 # Verify the SOCKS5 proxy works through WARP
790 echo -n " Verifying SOCKS5 proxy routes through WARP... "
791 if curl --socks5-hostname 127.0.0.1:$SOCKS_PROXY_PORT --max-time 10 -s https://www.cloudflare.com/cdn-cgi/
792 trace 2>/dev/null | grep -q "warp=on"; then
793     echo "ok (warp=on)."
794     echo " All TCP traffic now encrypted through Cloudflare WARP tunnel."
795 else
796     echo "FAILED (proxy not routing through WARP)."
797     echo " Disabling SOCKS5 proxy internet stays on direct connection."
798     disable_socks_proxy
799     echo " SOCKS proxy available at localhost:$SOCKS_PROXY_PORT for manual configuration."
800 fi
801 else
802     echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
803 fi
804
805 # 9. Verify host connectivity
806 if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
807     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
808         echo "Host is online on the Tailscale mesh."
809     else
810         echo "[Warning] Host is not responding on Tailscale."
811     fi
812 fi
813
814 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
815 echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."
816
817 # Prevent system from going to sleep while gateway is running
818 start_sleep_prevention
819 fi
820
821 SUCCESS_RUN=true
822
823 # Final DNS state summary
824 if [ -n "$ACTIVE_SERVICE" ]; then
825     FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>/dev/null | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr '
826 ' ' || true)
827     echo -e "\n"
828     echo -e "DNS STATE: $FINAL_DNS"
829     echo -e ""

```

```

829 fi
830
831 if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
832     if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
833         echo -e "\n[Warning] One or more blocklist URLs failed to download (404/Offline)."
```

The gateway gracefully fell back to yesterday's cached blocklist."

```

834         echo "
835         echo " (Run 'docker logs rule-compiler' later to investigate which link died.)"
836     fi
837 fi
838
839 echo -e "\n"
840 read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
841 if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
842     echo "Reversing gateway state..."
843     exec bash "$0"
844 fi
845
846 echo -e "\nYou can close this terminal window now."
847
848 exit 0
849 fi
850
851 #
852 # macOS TOGGLE jot / netstat / networksetup, launchd / pf.
853 #
854
855 # Enforce Cryptographic Script Integrity
856 if [ -f "scripts/crypto.sh" ]; then
857     if ! bash scripts/crypto.sh --verify; then
858         exit 1
859     fi
860 fi
861
862 # Define log file
863 LOG_FILE="$PWD/output.log"
864
865 # Rotate log file if it exceeds 1GB (1073741824 bytes).
866 # Sized large on purpose: with DEBUG_TRACE the log holds a full command-by-command
867 # trace effectively the entire compilation/warning history of the framework
868 # which is valuable to keep for post-mortems. Raised from 50MB so always-on
869 # tracing doesn't churn that history away prematurely.
870 if [ -f "$LOG_FILE" ]; then
871     log_size=$(stat -f%z "$LOG_FILE" 2>/dev/null || stat -c%s "$LOG_FILE" 2>/dev/null || echo 0)
872     if [ "$log_size" -gt 1073741824 ]; then
873         mv "$LOG_FILE" "${LOG_FILE}.old" 2>/dev/null || rm -f "$LOG_FILE"
874         echo "[$(date '+%Y-%m-%d %H:%M:%S')] Log file exceeded 1GB; rotated to output.log.old" > "$LOG_FILE"
875     fi
876 fi
877
878 # --- LIFECYCLE PHASE TRACKING + EXIT BREADCRUMB ---
879 # TOGGLE_PHASE is updated as the script moves through STARTUP/STOP/START so that
880 # if the process is killed (window closed, pkill, OOM) or exits with an error,
881 # the EXIT trap records the final exit code AND the phase it died in. Previously
882 # an interrupted START left output.log ending abruptly after the teardown with no
883 # indication of what failed this makes every exit traceable.
884 TOGGLE_PHASE="init"
885 _log_exit_breadcrumb() {
886     local rc=$?
887     echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh EXIT: code=$rc phase=$TOGGLE_PHASE (pid $$)" >> "$LOG_FILE"
888     2>/dev/null || true
889 }
890 trap '_log_exit_breadcrumb' EXIT
891
892 # --- DEBUG TRACE (opt-in via .env: DEBUG_TRACE=true) ---
893 # When enabled, mirror a full command-by-command xtrace to output.log for a
894 # complete post-mortem trail (e.g. a race during --restart while Wi-Fi is
895 # re-associating). Off by default. Read directly from .env here because
896 # common.sh (read_env_var) isn't sourced yet.
897 #
898 # BASH_XTRACEFD (send xtrace to its own fd, keeping the terminal clean) needs
899 # bash >= 4.1. macOS ships /bin/bash 3.2, so we branch:
900 #   - bash >= 4.1 (e.g. Linux): dedicate fd 9 to the log and point xtrace there;
901 #   - bash 3.2 (stock macOS): BASH_XTRACEFD is unsupported, so xtrace always goes
902 #     to stderr. We redirect stderr *straight* to the log file* with a plain
903 #     `exec 2>>"$LOG_FILE"`. This deliberately does NOT use a process
904 #     substitution (`2>>(tee)`): under `set -e`, stderr flowing through the
905 #     backgrounded `tee` proved to abort the script mid-START on 3.2 (a failed
906 #     write / SIGPIPE in a pipeline returned non-zero and tripped `set -e`, with
907 #     $LINENO mis-blaming the enclosing `fi` see devref 15.11.8-A/E). The
908 #     tradeoff: on 3.2 the terminal no longer shows live progress while tracing

```

```

909 # (everything is in output.log instead). Acceptable for an opt-in debug mode.
910 # We NEVER use the `exec {var}>>` dynamic-fd form (4.1+ only), which hard-fails on 3.2.
911 DEBUG_TRACE=$(grep -E '^DEBUG_TRACE=' .env 2>/dev/null | cut -d '=' -f2- | tr -d "\"" | tr '[:upper:]' '[:lower:]' | tr -d '[:space:]')
912 if [ "$DEBUG_TRACE" = "true" ]; then
913     echo "[$(date '+%Y-%m-%d %H:%M:%S')] DEBUG_TRACE enabled xtrace mirrored to output.log (bash ${BASH_VERSION}
914     [0]).${BASH_VERSION[1]})" >> "$LOG_FILE"
915     # Timestamped, location-aware xtrace prompt (works on 3.2 and 4.x).
916     export PS4='+ [$(SECONDS)s] ${BASH_SOURCE##*/}:${LINENO}:${FUNCNAME[0]:-main}: ' # subshell-free clock: a $
917     () in PS4 trips the bash 3.2 set -x/ERR heisenbug (devref 15.11.8)
918     if [ "${BASH_VERSION[0]}" -gt 4 ] || { [ "${BASH_VERSION[0]}" -eq 4 ] && [ "${BASH_VERSION[1]}" -ge 1 ];
919     }; then
920         exec 9>>"$LOG_FILE"
921         export BASH_XTRACEFD=9
922         set -x
923     else
924         # bash 3.2 (stock macOS): plain stderrfile redirect (no process substitution).
925         exec 2>>"$LOG_FILE"
926         set -x
927     fi
928 fi
929
930 # --- MAIN EXECUTION LOGGING ---
931 echo -e "\n===== " >> "$LOG_FILE"
932 echo "[$(date '+%Y-%m-%d %H:%M:%S')] toggle.sh invoked" >> "$LOG_FILE"
933 echo "===== " >> "$LOG_FILE"
934
935 # (Previously this script chmod 600'd .env at start to "unlock" it for docker compose,
936 # and chmod 000'd it on exit. That pattern broke docker compose on macOS/Colima when
937 # the umask produced 0600 the script removed its own ability to read .env before it could even proceed.)
938
939 # --- PROCEDURAL PORT GENERATION ---
940 # To prevent static port fingerprinting by malware, we randomize exposed ports
941 # on every boot and persist them to a hidden environment file.
942 PORTS_FILE="/tmp/nullexit-ports.env"
943 if [[ "$1" == "" || "$1" == "--restart" ]]; then
944
945     # Collision-avoidance function
946     GENERATED_PORTS=""
947     get_free_port() {
948         local port
949         while true; do
950             port=$(jot -r 1 10000 65000)
951             # 1. Prevent self-collision within this loop
952             if [[ "$GENERATED_PORTS" == *"$port"* ]]; then
953                 continue
954             fi
955             # 2. Prevent collision with ANY process on the Mac (root daemons, terminal apps, etc)
956             # We use netstat instead of lsof because netstat reads the kernel socket table
957             # directly and doesn't require sudo to see ports bound by other users/root.
958             if ! netstat -an -p tcp | awk '{print $4}' | grep -q "\.${port}$"; then
959                 echo "$port"
960                 return
961             fi
962         done
963     }
964
965     export TOR SOCKS_PORT=$(get_free_port)
966     GENERATED_PORTS="$GENERATED_PORTS $TOR SOCKS_PORT"
967
968     export TOR_CONTROL_PORT=$(get_free_port)
969     GENERATED_PORTS="$GENERATED_PORTS $TOR_CONTROL_PORT"
970
971     export SOCKS_PROXY_PORT=$(get_free_port)
972     GENERATED_PORTS="$GENERATED_PORTS $SOCKS_PROXY_PORT"
973
974     export DNS_PROXY_PORT=$(get_free_port)
975
976     echo "export TOR SOCKS_PORT=$TOR SOCKS_PORT" > "$PORTS_FILE"
977     echo "export TOR_CONTROL_PORT=$TOR_CONTROL_PORT" >> "$PORTS_FILE"
978     echo "export SOCKS_PROXY_PORT=$SOCKS_PROXY_PORT" >> "$PORTS_FILE"
979     echo "export DNS_PROXY_PORT=$DNS_PROXY_PORT" >> "$PORTS_FILE"
980     ADGUARD_PASS=$(grep -E '^ADGUARD_PASSWORD=' .env 2>/dev/null | cut -d '=' -f2- | tr -d "\"" )
981     echo "export TOR_PASSWORD=${ADGUARD_PASS:-nullexit}" >> "$PORTS_FILE"
982     elif [ -f "$PORTS_FILE" ]; then
983         # On STOP, source the existing ports so docker-compose down matches
984         source "$PORTS_FILE"
985     else
986         # Fallbacks if file is lost
987         export TOR SOCKS_PORT=9050
988         export TOR_CONTROL_PORT=9051

```

```

986 export SOCKS_PROXY_PORT=1080
987 export DNS_PROXY_PORT=5354
988 fi
989
990 # Start execution timer
991 TOGGLE_START_TIME=$SECONDS
992
993 if [[ "$1" == "--restart" ]]; then
994     echo "Executing Gateway Restart Sequence..."
995     # We must source common.sh here temporarily to check gateway state before we proceed
996     source "$SCRIPT_DIR/scripts/common.sh"
997     # Use persisted intent (marker) via gateway_state, which echoes a string (never
998     # a return code) to stay set -e-safe under set -x. See devref 15.11.8.
999     if [ "$(gateway_state)" = "running" ]; then
1000         echo "Gateway is currently running. Stopping it first..."
1001         bash "$0"
1002         echo "Gateway stopped. Now starting it..."
1003     else
1004         echo "Gateway is already stopped. Starting it..."
1005     fi
1006     bash "$0"
1007     exit 0
1008 fi
1009
1010 # Global flag to track if the script completed successfully
1011 SUCCESS_RUN=false
1012
1013 # Global variable to track the currently running background command PID
1014 CURRENT_BG_PID=""
1015
1016 # Source common bash functions
1017 source "$SCRIPT_DIR/scripts/common.sh"
1018
1019 # Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
1020 LOCK_FILE="/tmp/nullexit-toggle.lock"
1021 if [ -f "$LOCK_FILE" ]; then
1022     LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
1023     if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
1024         if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
1025             die "Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running."
1026         fi
1027     fi
1028 fi
1029 echo "$$" > "$LOCK_FILE"
1030
1031 # Helper function to run a GUI command as the logged-in console user
1032 # (Prevents permission failures if the script is run with sudo)
1033 run_gui_cmd() {
1034     local console_user
1035     console_user=$(stat -f '%Su' /dev/console 2>> output.log || echo "$SUDO_USER")
1036     if [ -z "$console_user" ] || [ "$console_user" = "root" ]; then
1037         console_user=$(logname 2>> output.log || echo "$USER")
1038     fi
1039
1040     if [ -n "$console_user" ] && [ "$console_user" != "root" ] && [ "$EUID" -eq 0 ]; then
1041         sudo -u "$console_user" "$@"
1042     else
1043         "$@"
1044     fi
1045 }
1046
1047 # Active-state marker: written at end of START path so external watchers
1048 # (see launchd/com.nullexit.wake-recovery.plist + scripts/watcher.sh +
1049 # devref 10.29) know the gateway is currently up. They only fire
1050 # recover.sh --post-wake when this file is present. Cleared at the top
1051 # of the STOP path and inside cleanup_handler on any error/signal path
1052 # so a half-broken gateway never re-triggers post-wake refreshes.
1053 write_gateway_active_marker() {
1054     date -u +%FT%TZ > /tmp/nullexit-gateway-active.marker
1055     # Set the watcher debounce timestamp to now to prevent a redundant post-startup recovery run.
1056     date +%s > /tmp/nullexit-watcher.last-recovery 2>/dev/null || true
1057 }
1058
1059 clear_gateway_active_marker() {
1060     rm -f /tmp/nullexit-gateway-active.marker
1061     rm -f "$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/TUNNEL_FAILED_CLOSED.marker"
1062 }
1063
1064 # Capture whether the gateway-active marker exists BEFORE the defensive clear
1065 # below wipes it. The STOP-vs-START decision (gateway_should_be_running) reads
1066 # this captured value persisted *intent* instead of a live Docker probe, so

```

```

1067 # a disable is instant and works even when Colima/Docker is down. Must be read
1068 # before clear_gateway_active_marker or it would always look "stopped".
1069 #
1070 # Written as default-assign + `if` WITHOUT an `else`: under `set -e` (with the
1071 # DEBUG_TRACE xtrace redirect active) an `if [ -f ]; then ; else ; fi` whose
1072 # condition is false was aborting the script here at init on macOS /bin/bash 3.2
1073 # the false `[ -f ]` status leaked through the compound. An `if` with no `else`
1074 # yields status 0 when the condition is false, so nothing leaks.
1075 GATEWAY_MARKER_AT_STARTUP="no"
1076 if [ -f "$MARKER_FILE" ]; then
1077     GATEWAY_MARKER_AT_STARTUP="yes"
1078 fi
1079
1080 # Defensive: clear any stale marker from a prior crashed/aborted run. If a
1081 # previous toggle.sh wrote the marker but was OOM-killed in the middle of the
1082 # START block (before the END-of-START write) or if the user rebooted mid
1083 # session this prevents the watcher from firing post-wake against a
1084 # non-running gateway on every wake event. Placed AFTER the function
1085 # definitions so bash resolves the call at parse time.
1086 clear_gateway_active_marker
1087
1088 PID_FILE="$PID_CAFFEINATE"
1089
1090 # Start macOS caffeinate to prevent sleep while gateway is running
1091 start_sleep_prevention() {
1092     if ! command -v caffeinate >/dev/null 2>&1; then
1093         echo " [Warning] caffeinate tool not found. Sleep prevention unavailable."
1094         return 1
1095     fi
1096
1097     # Stop any existing sleep prevention process first to avoid duplicates
1098     stop_pidfile_daemon "/tmp/nullexit-caffeinate.pid" "system sleep prevention"
1099
1100     echo " Enabling system sleep prevention (caffeinate)..."
1101     # Run caffeinate wrapped in a bash trap to prevent idle system sleep.
1102     # Critically, this traps macOS shutdown signals (SIGTERM) and automatically
1103     # flushes hijacked DNS back to normal right before the Mac powers off.
1104     # We do not use recover.sh here because it is a heavy recovery script (managing
1105     # containers, restarting tailscaled, etc.) which takes too long and gets
1106     # terminated by macOS before it can reset the DNS. Instead, we surgically and
1107     # instantly restore normal DNS and proxy settings here.
1108     nohup bash -c "
1109     trap '
1110         echo \"[Shutdown Trap] Reverting network settings...\" >> \"\$PWD/output.log\" 2>&1
1111         kill \$! 2>/dev/null || true
1112         sudo -n networksetup -setdnsservers \"\$ACTIVE_SERVICE\" 1.1.1.1 >> \"\$PWD/output.log\" 2>&1 || true
1113         sudo -n networksetup -setdnsservers \"\$ENO_SERVICE\" 1.1.1.1 >> \"\$PWD/output.log\" 2>&1 || true
1114         sudo -n networksetup -setsearchdomains \"\$ACTIVE_SERVICE\" \"Empty\" >> \"\$PWD/output.log\" 2>&1 || true
1115         sudo -n networksetup -setsearchdomains \"\$ENO_SERVICE\" \"Empty\" >> \"\$PWD/output.log\" 2>&1 || true
1116         sudo -n networksetup -setsocksfirewallproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 ||
1117     true
1118         sudo -n networksetup -setsocksfirewallproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1119         sudo -n networksetup -setwebproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1120         sudo -n networksetup -setwebproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1121         sudo -n networksetup -setsecurewebproxystate \"\$ACTIVE_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1122         sudo -n networksetup -setsecurewebproxystate \"\$ENO_SERVICE\" off >> \"\$PWD/output.log\" 2>&1 || true
1123         if command -v tailscale >/dev/null 2>&1; then
1124             tailscale up >> \"\$PWD/output.log\" 2>&1 || true
1125             tailscale set --accept-dns=false --exit-node= >> \"\$PWD/output.log\" 2>&1 || true
1126             TS_DOWN_ARGS=\"\"
1127             if [ \"\${KILL_SWITCH}\" = \"true\" ]; then TS_DOWN_ARGS=\"--accept-risk=lose-ssh\"; fi
1128             tailscale down \$TS_DOWN_ARGS >> \"\$PWD/output.log\" 2>&1 &
1129         fi
1130         sleep 1
1131         echo \"[Shutdown Trap] Cleanup complete.\" >> \"\$PWD/output.log\" 2>&1
1132         exit 0
1133     ' SIGTERM SIGINT SIGHUP
1134     caffeinate -i &
1135     wait \$!
1136     \" >> output.log 2>&1 &
1137     local caffe_pid=$!
1138     echo \"$caffe_pid\" > \"$PID_FILE"
1139     echo " Sleep prevention active (PID $caffe_pid). Your Mac won't sleep while the gateway is running."
1140 }
1141
1142
1143 DNS_WATCHER_PID_FILE="$PID_DNS_WATCHER"
1144 WARP_WATCHER_PID_FILE="/tmp/nullexit-warp-watcher.pid"
1145
1146

```

```

1147
1148 # Background WARP tunnel liveness monitor. Polls cdn-cgi/trace every 30s
1149 # while healthy, but accelerates to polling every 5s if a failure is detected.
1150 # Always logs state transitions to output.log. When warp=off persists for
1151 # WARP_FAIL_THRESHOLD consecutive polls (default 6 = 30s of downtime), the watcher
1152 # triggers a nuclear recovery: runs recover.sh to tear down the entire
1153 # gateway disconnecting Tailscale, stopping containers, resetting DNS,
1154 # and power-cycling Wi-Fi. Note: This minimizes exposure window, but
1155 # only the pf kill switch guarantees absolutely zero leakage on drop.
1156 #
1157 # The threshold (default 6) is statistically chosen: even at an
1158 # unrealistically high 10% false-positive rate per poll, P(6 consecutive
1159 # false positives) = 0.1^6 0.0001%. Real-world false-positive rates
1160 # are far lower, making 30s of continuous downtime overwhelmingly likely
1161 # to be a genuine outage.
1162 #
1163 # Configurable via .env: WARP_FAIL_THRESHOLD=3 (15s, more aggressive)
1164 start_warp_watcher() {
1165     stop_warp_watcher
1166     # Parse threshold from .env (same pattern as GATEWAY_MSS, GATEWAY_BYPASS_PING, etc.)
1167     local threshold
1168     threshold=$(read_env_var WARP_FAIL_THRESHOLD)
1169     if [ -z "$threshold" ]; then
1170         threshold="{WARP_FAIL_THRESHOLD:-6}"
1171     fi
1172     local trigger_file="/tmp/nullexit-warp-shutdown-triggered"
1173     rm -f "$trigger_file"
1174     echo " Starting background WARP Watcher (polling every 30s [accelerating to 5s on failure], shutdown after $
1175         {threshold} consecutive failures)..."
1176     nohup bash -c "
1177         trap 'exit 0' SIGTERM SIGINT SIGHUP
1178         last_state='on'
1179         consec_off=0
1180         threshold='$threshold'
1181         trigger_file='$trigger_file'
1182         out_log='$PWD/output.log'
1183         recover_bin='$PWD/recover.sh'
1184         inhibit_file='/tmp/nullexit-warp-inhibit.marker'
1185         while true; do
1186             # Inhibit check
1187             # recover.sh --post-wake writes this marker while force-recreating the
1188             # warp container. During that window, the container is intentionally
1189             # down don't count it as a failure or we'll fire nuclear recover.sh
1190             # and kill a gateway that was in the process of healing itself.
1191             if [ -f "\$inhibit_file" ]; then
1192                 consec_off=0
1193                 sleep 5
1194                 continue
1195             fi
1196             state='off'
1197             if docker compose exec -T warp wget -qO- --timeout=5 \
1198                 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
1199                 | grep -q 'warp=on'; then
1200                 state='on'
1201             fi
1202             # State-transition logging
1203             if [ "\$state" != "\$last_state" ]; then
1204                 if [ "\$state" = 'off' ]; then
1205                     echo "\[$(date -u +%FT%TZ)] WARP DOWN cdn-cgi/trace reports warp=off\" >> \"\$out_log\"
1206                 fi
1207                 last_state="\$state"
1208             fi
1209             # Consecutive-failure tracking + auto-shutdown
1210             if [ "\$state" = 'off' ]; then
1211                 consec_off=$((consec_off + 1))
1212                 if [ "\$consec_off" -ge "\$threshold" ] && [ ! -f "\$trigger_file" ]; then
1213                     touch "\$trigger_file"
1214                     echo "\[$(date -u +%FT%TZ)] WARP SHUTDOWN \$consec_off consecutive failures (threshold=\$threshold)
1215                     . Running recover.sh to kill the gateway and restore normal internet. Your IP is NO LONGER Cloudflare.\" >>
1216                     \"\$out_log\"
1217                     # Nuclear recovery: tear down the entire gateway.
1218                     # recover.sh resets DNS 1.1.1.1, disconnects Tailscale,
1219                     # stops Docker containers, flushes routes, power-cycles Wi-Fi.
1220                     bash "\$recover_bin" --auto >> \"\$out_log\" 2>&1 || true
1221                     echo "\[$(date -u +%FT%TZ)] WARP SHUTDOWN recover.sh completed, watcher exiting. Your IP is now
1222                     your real ISP IP.\" >> \"\$out_log\"
1223                     exit 0
1224                 fi

```



```

1224     else
1225         if [ \("${consec_off}" -gt 0 ); then
1226             echo \("${date -u +%FT%TZ}") WARP RECOVERED warp=on after \("${consec_off}" consecutive off readings (
threshold was \("${threshold}") >> \("${out_log}"
1227             fi
1228             consec_off=0
1229             fi
1230
1231             if [ \("${state}" = "off" ); then
1232                 sleep 5
1233             else
1234                 sleep 30
1235             fi
1236             done
1237             " >> output.log 2>&1 &
1238             echo $! > "$WARP_WATCHER_PID_FILE"
1239         }
1240
1241         stop_warp_watcher() {
1242             if [ -f "$WARP_WATCHER_PID_FILE" ]; then
1243                 local wp
1244                 wp=$(cat "$WARP_WATCHER_PID_FILE")
1245                 if [ -n "$wp" ] && kill -0 "$wp" 2>/dev/null; then
1246                     echo " Stopping background WARP Watcher (PID $wp)..."
1247                     kill "$wp" 2>/dev/null || true
1248                 fi
1249                 rm -f "$WARP_WATCHER_PID_FILE"
1250             fi
1251         }
1252
1253
1254
1255
1256 # Cleanup handler to restore DNS to 1.1.1.1 on error or user interrupt (Ctrl+C / SIGTERM / SIGHUP)
1257 cleanup_handler() {
1258     local exit_code=$?
1259     local trigger_type="$1"
1260     local line_no="$2"
1261
1262     if [ "$SUCCESS_RUN" = "false" ]; then
1263         echo -e "\n[Self-Correction] Script interrupted or failed ($trigger_type). Restoring host DNS to 1.1.1.1..."
1264         "
1265         if [ -n "$ACTIVE_SERVICE" ]; then
1266             reset_dns
1267         fi
1268
1269         if [ -n "$CURRENT_BG_PID" ]; then
1270             echo "Terminating active command (PID $CURRENT_BG_PID)..."
1271             kill -15 "$CURRENT_BG_PID" 2>> output.log || true
1272         fi
1273
1274         # Disconnect host Tailscale and close the GUI application on failure
1275         disconnect_tailscale_host
1276
1277         # Stop local DNS proxy if running
1278         stop_local_dns_proxy
1279
1280         # Stop sleep prevention
1281         stop_pidfile_daemon "/tmp/nullexit-cafeinate.pid" "system sleep prevention"
1282         stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
1283         stop_warp_watcher
1284         clear_gateway_active_marker
1285
1286         # Capture warp logs on failure for debugging before teardown
1287         if [ "$trigger_type" = "ERR" ]; then
1288             echo -e "\n--- WARP FAILURE LOGS ---" >> output.log
1289             docker compose logs warp --tail=100 >> output.log 2>&1 || true
1290             echo "-----" >> output.log
1291         fi
1292
1293         # Best-effort container teardown on failure
1294         docker compose down --remove-orphans -t 30 2>> output.log || true
1295
1296         # Nuke leftover network state (proxies, routes, DNS cache, Wi-Fi)
1297         if [ -n "$ACTIVE_SERVICE" ]; then
1298             cleanup_network_state
1299         fi
1300
1301         if [ "$trigger_type" = "ERR" ]; then
1302

```

```

1303     echo -e "\n=====
1304     echo "ERROR: Script failed at line $line_no with exit code $exit_code."
1305     echo "=====
1306     echo "Troubleshooting steps:"
1307     echo "1. Run 'colima status' to verify the VM is active."
1308     echo "2. Run 'docker compose ps' to check container health."
1309     echo "3. Run 'docker compose logs' to view application logs."
1310     echo "4. Run 'tailscale status' to check host Tailscale status."
1311     echo "=====
1312 fi
1313 rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
1314 fi
1315 }
1316
1317 trap 'cleanup_handler ERR $LINENO' ERR
1318 trap 'cleanup_handler INT' INT
1319 trap 'cleanup_handler TERM' TERM
1320 trap 'cleanup_handler HUP' HUP
1321
1322 # NOPASSWD sudo
1323 # The user has NOPASSWD in /etc/sudoers.d/nullexit for the specific commands
1324 # needed (networksetup, pfctl, route, ifconfig, dscacheutil, killall, pkill,
1325 # and python3 -I scripts/dns-proxy.py for the fallback DNS proxy).
1326 # All privileged calls below use `sudo -n` which will run without prompting.
1327 # No credential caching loop is needed.
1328
1329 # Add Homebrew and standard paths
1330 setup_standard_path
1331
1332 # Check for local DNS proxy tools (used when tailnet data plane is unavailable)
1333 # Uses a Python DNS proxy (handles TCP wire format correctly).
1334 # Python3 is built into macOS no external dependencies.
1335 DNS_PROXY_BIN=""
1336 if command -v python3 >> output.log 2>&1; then
1337     DNS_PROXY_BIN="python3 -I"
1338 elif command -v python >> output.log 2>&1; then
1339     DNS_PROXY_BIN="python -I"
1340 fi
1341
1342 # Path to the DNS proxy script (sibling of this script)
1343 DNS_PROXY_SCRIPT="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/scripts/dns-proxy.py"
1344
1345 # Resolve Tailscale CLI on host. We rely on the standalone Homebrew formula
1346 # (`brew install tailscale` + `brew services start tailscale`) NOT the
1347 # Tailscale.app GUI so there is no .app bundle or Network Extension sandbox
1348 # to launch, no menu-bar click to perform, and no scutil Network Service to
1349 # start manually. tailscaled runs as a per-user LaunchAgent (or LaunchDaemon
1350 # if you ran the install with sudo) and the control socket is always available.
1351 TS_BIN=""
1352 if command -v tailscale >> output.log 2>&1; then
1353     TS_BIN="tailscale"
1354 fi
1355
1356 # Baseline: disable Tailscale DNS management at the daemon level.
1357 # `tailscale set` writes a persistent preference. However, `tailscale up --reset`
1358 # (used in steps 7 and 9) RESETS all unspecified flags to defaults, which would
1359 # re-enable `--accept-dns=true` and nullify this `set`.
1360 # Therefore, EVERY `tailscale up --reset` call in this file MUST also pass
1361 # `--accept-dns=false` explicitly. See steps 7 and 9 for the fix.
1362 #
1363 # This initial `set` is still useful as a baseline for non-reset operations
1364 # (e.g. if the user runs `tailscale up` manually without --reset) and covers
1365 # the brief window between this line and the first `--reset` call.
1366 #
1367 # Runs once at script start while tailscaled is still connected (if it was
1368 # left running from a previous toggle), so the pref takes effect before any
1369 # disconnection or reconnection logic.
1370 if [ -n "$TS_BIN" ]; then
1371     $TS_BIN set --accept-dns=false >> output.log 2>&1 || true
1372 fi
1373
1374 # Disconnect host Tailscale from the mesh. With the standalone daemon, this is the
1375 # *entire* teardown no scutil tunnel stop, no Tailscale.app GUI quit, no pkill.
1376 # tailscaled keeps running (as a LaunchAgent / LaunchDaemon), which is intentional:
1377 # the next `tailscale up` then reconnects in seconds rather than waiting for a
1378 # slow .app launch + Network Extension activation. The `tailscale down` call also
1379 # retains the user's auth state, so we don't force a browser re-login every cycle.
1380 #
1381 # Defined early (before the if/else branches and referenced by cleanup_handler's
1382 # ERR trap) so that any failure before the main logic can still tear down Tailscale.
1383

```

```

1384
1385 # Wait for Network Readiness
1386 # On --restart, the gateway tears down while the host network interface may
1387 # still be reconnecting (Wi-Fi, Ethernet, USB tethering, etc.). If we proceed
1388 # too early, get_active_service returns nothing, routing targets the wrong
1389 # interface, and DNS/WARP setup silently fails.
1390 #
1391 # We use `route get default` to detect a valid default gateway rather than
1392 # checking a specific interface (e.g. en0) this generalizes across all
1393 # connection types: Wi-Fi, Ethernet, USB-C adapters, iPhone tethering, etc.
1394 # The moment the OS has any routable default gateway, we proceed.
1395 _net_wait_secs=0
1396 _net_timeout=60
1397 echo "Waiting for network connectivity before starting..."
1398 while true; do
1399     if route get default 2>/dev/null | grep -q 'gateway: '; then
1400         _gw=$(route get default 2>/dev/null | awk '/gateway:/{print $2}')
1401         echo "    Network ready (default gateway: $_gw)."
1402         break
1403     fi
1404     if [ "$_net_wait_secs" -ge "$_net_timeout" ]; then
1405         echo "    [Warning] No default gateway after ${_net_timeout}s proceeding anyway."
1406         break
1407     fi
1408     sleep 2
1409     _net_wait_secs=$(( _net_wait_secs + 2 ))
1410 done
1411
1412 ACTIVE_SERVICE=$(get_active_service)
1413
1414 # Resolve the service name for en0 (usually "Wi-Fi") macOS scutil DNS resolver
1415 # is commonly scoped to en0, so per-service DNS changes on other interfaces are
1416 # ignored unless en0's service is also updated.
1417 ENO_SERVICE=$(get_en0_service)
1418
1419 # Helper to add host-side bypass routes for the WARP WireGuard endpoints.
1420 # This prevents an infinite recursive routing loop (tunnel loop) where
1421 # the container's WireGuard packets exit the VM, reach the host, match the
1422 # default route (the exit node), and get routed back into the tunnel.
1423 # We route via the gateway IP (not interface) to ensure packets are routed
1424 # properly even when the host's default route changes to the Tailscale interface.
1425
1426 # Helper to manually override the host's default route to point through the
1427 # Tailscale utun* interface. Standalone tailscaled on macOS (Homebrew version)
1428 # lacks the platform-specific integration to override the default gateway
1429 # automatically when --exit-node is used.
1430 setup_exit_node_routing() {
1431     local ts_iface
1432     ts_iface=$(ifconfig | awk '/^[a-z0-9]+:/{iface=$1} /inet 100\./{print iface; exit}' | tr -d ':')
1433
1434     if [ -n "$ts_iface" ]; then
1435         echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
1436         local host_mtu
1437         host_mtu=$(read_env_var HOST_MTU)
1438         host_mtu=${host_mtu:-1200}
1439
1440         # Lower the host Tailscale MTU (defaults to 1200) to prevent fragmentation when passing
1441         # through the container's WARP tunnel (which also has an MTU of 1280).
1442         sudo -n ifconfig "$ts_iface" mtu "$host_mtu" >> output.log 2>&1 || true
1443
1444         # DO NOT delete the physical default route! It crashes Colima.
1445         # Use the /1 trick to mathematically override the default route.
1446         sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
1447         sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
1448         sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
1449         sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
1450
1451         if netstat -nr | grep -E -q "^(0\.\0\.\0/1|0/1)[[:space:]]+.*$ts_iface"; then
1452             echo "    [] Default route successfully overridden via $ts_iface."
1453         else
1454             echo "    [!] Failed to verify exit node routing on $ts_iface."
1455         fi
1456     else
1457         echo "[Warning] Could not detect Tailscale utun interface for host routing."
1458     fi
1459 }
1460
1461 # Helper to nuke leftover network state after teardown (proxy settings, routing
1462 # table, DNS cache, Wi-Fi). Prevents the "had to write a whole recovery script"
1463 # problem where stale state kills internet after the gateway stops.
1464 cleanup_network_state() {

```

```

1465 echo -e "\nCleaning up network state (proxies, routes, DNS cache, Wi-Fi)..."
1466
1467 # 1. Disable any leftover proxy settings (needs root on many macOS versions)
1468 for svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
1469     disable_all_proxies "$svc"
1470 done
1471 echo "    Proxies disabled."
1472
1473 # 2. Flush DNS cache
1474 if command -v dscacheutil >> output.log 2>&1; then
1475     sudo -n dscacheutil -flushcache 2>> output.log || true
1476 fi
1477 sudo -n killall -HUP mDNSResponder 2>> output.log || true
1478 echo "    DNS cache flushed."
1479
1480 # 3. Flush stale routing table entries
1481 if command -v route >> output.log 2>&1; then
1482     sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
1483     sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
1484 fi
1485 remove_warp_bypass_routes
1486 disable_killswitch
1487 echo "    Routing table and firewall flushed."
1488
1489 # 4. Power-cycle Wi-Fi to clear any lingering interface state
1490 WIFI_PORT=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print $2
1491 }')
1492 if [ -n "$WIFI_PORT" ]; then
1493     sudo -n networksetup -setairportpower "$WIFI_PORT" off 2>> output.log || true
1494     sleep 2
1495     sudo -n networksetup -setairportpower "$WIFI_PORT" on 2>> output.log || true
1496     echo "    Wi-Fi power-cycled (interface $WIFI_PORT)."
1497     sleep 3
1498 else
1499     echo "    Wi-Fi interface not detected; bouncing en0..."
1500     sudo -n ifconfig en0 down 2>> output.log || true
1501     sudo -n ifconfig en0 up 2>> output.log || true
1502 fi
1503 # Force restart sharing services to prevent AirDrop / AirPlay discovery freezes
1504 # after network interface/DNS changes.
1505 echo "    Resetting macOS sharing services (AirDrop/AirPlay)..."
1506 reset_sharing_services
1507
1508 echo "Network state cleanup complete."
1509 }
1510
1511 SOCKS_PROXY_PORT=${SOCKS_PROXY_PORT}
1512
1513 # Ensure DNS is set to 1.1.1.1 immediately at script start to prevent any DNS deadlocks/hangs
1514 # during status checks, container startup, VM boot, or teardown. We *also* clear any leftover
1515 # `ts.net` search domain from a previous `tailscale up --accept-dns=true` run, otherwise macOS
1516 # would prepend it to every lookup and most public DNS queries would resolve to NXDOMAIN.
1517
1518
1519 echo "Initializing DNS to 1.1.1.1 to ensure reliable internet access..."
1520 reset_dns
1521
1522
1523
1524
1525 echo "Checking Gateway Status..."
1526 HIJACK_HOST=$(read_env_var GATEWAY_HIJACK_HOST | tr '[:upper:]' '[:lower:]')
1527
1528 # Decide STOP vs START from persisted intent (marker), not a live Docker probe.
1529 # This makes *disable* instant and robust even when Colima/Docker is down.
1530 #
1531 # CRITICAL (set -e + set -x on bash 3.2): gateway_state ECHOES "running"/"stopped"
1532 # and always returns 0, so we branch on the STRING never on a function's exit
1533 # code. A helper that `return 1`'s trips a spurious `set -e` abort under `set -x`
1534 # (DEBUG_TRACE=true) even when guarded, hence the echo contract. See devref 15.11.8.
1535 if [ "$(gateway_state)" = "running" ]; then
1536     TOGGLE_PHASE="stop"
1537     echo -e "\n=====
1538 echo "Gateway is RUNNING. Stopping it now..."
1539 echo -e "=====
1540
1541 # 1. Disconnect host Tailscale cleanly first before stopping the exit-node container
1542 if [[ "$HIJACK_HOST" != "false" ]]; then
1543     disconnect_tailscale_host
1544     echo ""

```

```

1545 fi
1546
1547 echo "Stopping Docker containers..."
1548 # `|| true`: a disable must always complete its teardown even if the Docker
1549 # daemon / Colima VM is already down (compose down would exit non-zero and,
1550 # under set -e, abort the STOP). The START path intentionally does NOT guard
1551 # its `docker compose up` a failed start SHOULD trigger cleanup.
1552 run_logged docker compose down --remove-orphans -t 30 || true
1553
1554 # Only stop Colima if the user explicitly opted in (false by default) to prevent breaking their other Docker
1555 # dev projects
1556 STOP_COLIMA=$(read_env_var STOP_COLIMA_ON_EXIT | tr '[:upper:]' '[:lower:]')
1557 if [[ "$STOP_COLIMA" == "true" ]]; then
1558     echo -e "\nStopping Colima VM to free up host RAM and battery..."
1559     run_with_timeout 30 colima stop >> output.log 2>&1 || echo "Warning: Failed to stop Colima gracefully."
1560 else
1561     echo -e "\nLeaving Colima running. (Set STOP_COLIMA_ON_EXIT=true in .env to change this)"
1562 fi
1563
1564 if [[ "$HIJACK_HOST" != "false" ]]; then
1565     # The host's DNS was hijacked to the gateway IP during ENABLE; now that the
1566     # gateway is down, restore DNS to 1.1.1.1 immediately so subsequent lookups
1567     # don't stall for the macOS DNS timeout before the 1.1.1.1 fallback engages.
1568     echo -e "\nRestoring host DNS to 1.1.1.1 (gateway is gone)... "
1569     reset_dns
1570
1571     # 3. Stop local DNS proxy if running
1572     stop_local_dns_proxy
1573
1574     # Stop background daemons
1575     stop_pidfile_daemon "/tmp/nullexit-cafeinate.pid" "system sleep prevention"
1576     stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
1577     stop_warp_watcher
1578     clear_gateway_active_marker
1579
1580     # 4. Nuke leftover network state so internet actually works after teardown
1581     cleanup_network_state
1582 fi
1583
1584 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
1585 echo -e "\nGateway has been successfully STOPPED in ${ELAPSED} seconds."
1586 else
1587     TOGGLE_PHASE="start"
1588     START_GATEWAY=true
1589     echo "[$(date +%Y-%m-%d %H:%M:%S)] Action: STARTUP GATEWAY" >> "$LOG_FILE"
1590     echo -e "\n\n=====
1591     echo "Gateway is STOPPED. Starting it now..."
1592     echo -e "=====
1593
1594     # 1. Prevent host exit-node deadlock during VM / Container startup
1595     if [[ "$HIJACK_HOST" != "false" ]]; then
1596         disconnect_tailscale_host
1597     fi
1598
1599     # 2. Wait for physical DHCP lease to settle (crucial for restarts/wake-up/roaming)
1600     wait_for_dhcp_settle
1601
1602     # 3. Boot Colima VM if it is not already running
1603     echo -e "\nChecking Colima VM status..."
1604     if ! run_with_timeout 15 colima status >> output.log 2>&1; then
1605         # Colima networking is gated on the SAME LAN-P2P signal routing-fix uses
1606         # (.lan_p2p_detected overrides TAILSCALE_ALLOW_LAN_P2P in .env). A LAN-reachable
1607         # VM (--network-address + bridged/shared) needs the socket_vmnet helper and is
1608         # only worthwhile on a trusted home network, so we request it only when P2P is
1609         # allowed; otherwise AP-isolated / enterprise, the common case we boot plain
1610         # user-mode networking. Either way host<->container P2P is UNAFFECTED: it rides
1611         # routing-fix's Docker-gateway / .host_ips RETURN rules, not the VM network mode
1612         # (see devref 15.3.5). colima_start_until_ready returns as soon as Docker answers
1613         # (~15s) instead of blocking on colima's ~2min post-provision hang (see 15.8.7);
1614         # that hang is colima-internal and mode-independent (it also occurs in user mode).
1615         allow_p2p=$(read_env_var "TAILSCALE_ALLOW_LAN_P2P")
1616         [ -z "$allow_p2p" ] && allow_p2p="false"
1617         if [ -f "$SCRIPT_DIR/.lan_p2p_detected" ]; then
1618             allow_p2p=$(tr -d '[:space:]' < "$SCRIPT_DIR/.lan_p2p_detected" 2>/dev/null || echo "false")
1619         fi
1620
1621         if [ "$allow_p2p" = "true" ]; then
1622             colima_network_mode="bridged"
1623             if [[ "$OSTYPE" == "darwin*" ]]; then
1624                 security_type=$(system_profiler SPAirPortDataType 2>/dev/null | awk '/Current Network Information:/{
1625                     found=1} found && /Security:/{print $NF; exit}')

```

```

1624         if echo "$security_type" | grep -qi "Enterprise"; then
1625             echo -e "\n [!] WPA2-Enterprise Wi-Fi detected! Bridged mode will cause MAC-spoofing
deauthentication."
1626             echo "             Falling back to '--network-mode shared' for Colima."
1627             colima_network_mode="shared"
1628         fi
1629         fi
1630         echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P on: --network-address
$colima_network_mode)..."
1631         colima_start_until_ready --memory 0.6 --vm-type vz --network-address --network-mode "$colima_network_mode
" \
1632         || echo " [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry
)."
1633     else
1634         colima_network_mode="user"
1635         echo "Colima is not running. Starting Colima (600MB RAM, vz VM, LAN P2P off: fast user-mode networking)
..."
1636         colima_start_until_ready --memory 0.6 --vm-type vz \
1637         || echo " [!] Colima Docker not confirmed ready; continuing (downstream docker-ps auto-heal will retry
)."
1638         fi
1639         echo "$colima_network_mode" > /tmp/nullexit_colima_mode.txt
1640     else
1641         echo "Colima is already running."
1642     fi
1643 fi
1644 # 3b. Auto-detect and fix Docker Subnet Collisions (e.g. if the user travels to a 172.17.x.x Wi-Fi)
1645 if [ -f "scripts/fix-docker-bridge-collision.sh" ]; then
1646     if bash scripts/fix-docker-bridge-collision.sh --dry | grep -q "collision detected"; then
1647         echo -e "\n[!] WARNING: Docker Subnet Collision Detected with your current Wi-Fi!"
1648         echo -e "             Auto-fixing Colima config to prevent internet jitter..."
1649         bash scripts/fix-docker-bridge-collision.sh --skip-toggle || true
1650     fi
1651 fi
1652 # 3c. Configure swap file inside the VM to prevent OOM on low-memory limits
1653 # We also set vm.swappiness=10 so the Linux kernel strictly prefers physical RAM
1654 # and avoids unnecessarily wearing out the SSD with proactive background swapping.
1655 if ! run_with_timeout 15 colima ssh -- grep -q 'swapfile' /proc/swaps >> output.log 2>&1; then
1656     echo "Configuring 400MB swap file inside the VM to prevent OOM..."
1657     run_with_timeout 30 colima ssh -- sudo sh -c "if [ ! -f /swapfile ]; then dd if=/dev/zero of=/swapfile bs=1
M count=400 status=none && mkswap /swapfile; fi && swapon /swapfile && sysctl vm.swappiness=10" >> output.
log 2>&1 || echo "Warning: Failed to enable swap file inside the VM."
1659 fi
1660 # 4. Clean up corrupted AdGuardHome configurations
1661 if [ -f "adguard/conf/AdGuardHome.yaml" ] && [ ! -s "adguard/conf/AdGuardHome.yaml" ]; then
1662     rm -f "adguard/conf/AdGuardHome.yaml"
1663     echo "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop."
1664 fi
1665 # 4b. Auto-heal stale Docker socket from hard reboots
1666 if ! run_with_timeout 15 docker ps >/dev/null 2>&1; then
1667     echo "Docker daemon is unresponsive (likely a stale socket from a crash). Auto-healing Colima..."
1668     run_with_timeout 120 colima restart >> output.log 2>&1
1669 fi
1670 # 4c. Inject TCP MSS clamp from .env into post-rules.txt before starting
1671 if [ -f .env ] && grep -q "^GATEWAY_MSS=" .env; then
1672     GATEWAY_MSS=$(read_env_var GATEWAY_MSS)
1673     if [[ "$GATEWAY_MSS" =~ ^[0-9]+$ ]]; then
1674         # macOS sed syntax
1675         sed -i '' "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null || \
1676         # Linux sed fallback
1677         sed -i "s/--set-mss [0-9]*/--set-mss ${GATEWAY_MSS}/g" post-rules.txt 2>/dev/null
1678     fi
1679 fi
1680 # 5. Start compose services
1681 write_host_ips
1682 # Procedurally generated ports have already been exported at the top of the script.
1683 echo " [Security] Procedurally randomized proxy ports generated to evade fingerprinting."
1684 # --- HONEY-PORT TRIPWIRE ---
1685 # Generate a random trap port. If any malware does a sequential port scan on localhost,
1686 # it will inevitably hit this port. When it does, we instantly fire a macOS alert.
1687 # Reap the honey-port from any previous toggle before arming a new one, so the
1688 # `nc -l` tripwire never accumulates one idle listener per toggle-on (15.12.20).
1689 pkill -f "nc -l localhost" 2>/dev/null || true
1690 HONEY_PORT=$(get_free_port)

```



```

1697 (
1698     if nc -l localhost "$HONEY_PORT" >/dev/null 2>&1; then
1699         echo "[$(date +%Y-%m-%d %H:%M:%S)] [CRITICAL SECURITY ALERT] PORT SCAN DETECTED! An unknown application
1700             just pinged the local Honey-Port ($HONEY_PORT)!" >> "$LOG_FILE"
1701         osascript -e 'display notification "An unknown application is aggressively scanning your local ports.
1702             Malware port-scan suspected." with title "nullexit OPSEC Alert" sound name "Basso" 2>/dev/null || true
1703     fi
1704 ) &
1705 disown %1 2>/dev/null || true
1706 echo " [Security] Honey-Port Tripwire armed on random port to detect malware port-scans."
1707
1708 echo -e "\nStarting Docker containers..."
1709
1710 # --- Ad-block rule compilation: hash-gated one-off, off the critical path (15.8.8) ---
1711 # The rule-compiler (a `compile`-profiled container, the SOLE writer of adguard/work)
1712 # re-dedupes ~340k rules inside the 600MB VM (~28s) far too slow to run every toggle.
1713 # It is gated on a hash of the lists YOU control (black_list.txt + white_list.txt):
1714 # recompile only when they change, when compiled_rules.txt / ip_blocklist.ipset are
1715 # missing, or when they exceed RULES_MAX_AGE_DAYS (default 7, so the remote blocklists
1716 # still refresh occasionally). Otherwise AdGuard/routing-fix boot on the existing
1717 # artifacts and we skip the ~28s they no longer `depends_on` the rule-compiler
1718 # container (see docker-compose.yml). A compile failure is non-fatal: we keep the
1719 # previous compiled rules rather than bricking the tunnel (ad-block != the kill-switch).
1720 RULES_HASH_FILE="$SCRIPT_DIR/.rules_hash"
1721 rules_hash=$(compute_rules_hash)
1722 stored_rules_hash=""
1723 [ -f "$RULES_HASH_FILE" ] && stored_rules_hash=$(cat "$RULES_HASH_FILE" 2>/dev/null)
1724 rules_max_age_days=$(read_env_var RULES_MAX_AGE_DAYS)
1725 [ -z "$rules_max_age_days" ] && rules_max_age_days=7
1726 need_compile=false
1727 if [ ! -f adguard/work/userfilters/compiled_rules.txt ] || [ ! -f adguard/work/userfilters/ip_blocklist.ipset
1728 ]; then
1729     need_compile=true
1730 elif [ -z "$rules_hash" ] || [ "$rules_hash" != "$stored_rules_hash" ]; then
1731     need_compile=true
1732 elif [ -n "$(find adguard/work/userfilters/compiled_rules.txt -mtime +$rules_max_age_days 2>/dev/null)" ]; then
1733     need_compile=true
1734 fi
1735 if [ "$need_compile" = "true" ]; then
1736     echo " Ad-block lists changed (or artifacts missing/stale) compiling rules (one-off, ~28s)..."
1737     if run_logged docker compose run --build --rm rule-compiler >> output.log 2>&1; then
1738         if [ -n "$rules_hash" ]; then echo "$rules_hash" > "$RULES_HASH_FILE"; fi
1739     else
1740         echo " [!] Rule compile failed continuing with existing compiled ad-block rules."
1741     fi
1742 else
1743     echo " Ad-block lists unchanged reusing compiled rules (skipping the ~28s recompile)."
1744 fi
1745
1746 # --- Gate `--build` for the long-running images (routing-fix, tor) ---
1747 # BuildKit re-checking these on the 600MB VM costs ~30s even when every layer is
1748 # CACHED (15.8.8). Only pass --build when the hashed inputs change or an image is
1749 # missing; otherwise plain `up -d` reuses cached images. `up` starts everything
1750 # except the `compile`-profiled rule-compiler.
1751 BUILD_HASH_FILE="$SCRIPT_DIR/.build_hash"
1752 build_inputs_hash=$(compute_build_inputs_hash)
1753 stored_build_hash=""
1754 [ -f "$BUILD_HASH_FILE" ] && stored_build_hash=$(cat "$BUILD_HASH_FILE" 2>/dev/null)
1755 local_images_present=true
1756 for _img in nullexit-routing-fix:v1.0.0 nullexit-tor:v1.0.0; do
1757     docker image inspect "$_img" >> output.log 2>&1 || local_images_present=false
1758 done
1759 if [ -n "$build_inputs_hash" ] && [ "$build_inputs_hash" = "$stored_build_hash" ] && [ "$local_images_present"
1760     = "true" ]; then
1761     echo " Build inputs unchanged skipping image build (reusing cached images)."
1762     run_logged docker compose up -d
1763 else
1764     echo " Build inputs changed or a local image is missing building images..."
1765     run_logged docker compose up -d --build
1766     # Record the hash only after a successful build (set -e aborts above on failure).
1767     if [ -n "$build_inputs_hash" ]; then echo "$build_inputs_hash" > "$BUILD_HASH_FILE"; fi
1768 fi
1769
1770 RULE_COUNT=$(grep "! Total Block Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -F': ' '{print $2}' || echo "0")
1771 NATIVE_COUNT=$(grep "! Native AdGuard Rules:" adguard/work/userfilters/compiled_rules.txt 2>/dev/null | awk -F': ' '{print $2}' || echo "0")
1772
1773 if [ "$RULE_COUNT" != "0" ] && [ "$RULE_COUNT" != "" ]; then
1774     if [ "$NATIVE_COUNT" != "0" ] && [ "$NATIVE_COUNT" != "" ]; then

```

```

1771     TOTAL_COUNT=$((RULE_COUNT + NATIVE_COUNT))
1772     # Format with commas for readability (e.g. 441,578)
1773     if command -v printf &> /dev/null; then
1774         FORMATTED_RULE=$(printf "%'d" "$RULE_COUNT")
1775         FORMATTED_TOTAL=$(printf "%'d" "$TOTAL_COUNT")
1776         echo "   DNS: Compiled $FORMATTED_RULE unique custom rules (Total active in AdGuard: ~$FORMATTED_TOTAL
rules)."
1777     else
1778         echo "   DNS: Compiled $RULE_COUNT unique custom rules (Total active in AdGuard: ~$TOTAL_COUNT rules)."
1779     fi
1780 else
1781     echo "   DNS: Compiled and loaded $RULE_COUNT optimized rules."
1782 fi
1783 fi
1784
1785 # Report IP blocklist compilation results
1786 IP_COUNT=$(grep "^# Entries:" adguard/work/userfilters/ip_blocklist.ipset 2>/dev/null | awk -F': ' '{print $2
}' || echo "0")
1787 if [ "$IP_COUNT" != "0" ] && [ "$IP_COUNT" != "" ]; then
1788     if command -v printf &> /dev/null; then
1789         FORMATTED_IP=$(printf "%'d" "$IP_COUNT")
1790         echo "   IP:   Loaded $FORMATTED_IP threat intelligence IPs/CIDRs into kernel firewall."
1791     else
1792         echo "   IP:   Loaded $IP_COUNT threat intelligence IPs/CIDRs into kernel firewall."
1793     fi
1794 fi
1795
1796 # 6. Wait for the gateway container's Tailscale connection to be ready
1797 BYPASS_PING=$(read_env_var GATEWAY_BYPASS_PING | tr '[:upper:]' '[:lower:]')
1798 USE_EXIT_NODE=$(read_env_var GATEWAY_USE_EXIT_NODE | tr '[:upper:]' '[:lower:]')
1799
1800 echo "Waiting for gateway container's Tailscale to connect to the tailnet..."
1801 echo " (This can take 30-60s  Tailscale goes through: NoState Starting Running Online)"
1802 connected=false
1803 consecutive_failures=0
1804 for i in {1..60}; do
1805     # Run tailscale status and capture its output so the user can see what's happening.
1806     # Previously this was hidden behind 2>> output.log, making it impossible to debug hangs.
1807     ts_output=$(run_with_timeout 3 docker compose exec -T tailscale tailscale status 2>&1 || true)
1808
1809     if echo "$ts_output" | grep -q "offers exit node"; then
1810         connected=true
1811         break
1812     fi
1813
1814     # Track consecutive failures to detect a stuck state.
1815     # If we see 40+ consecutive 'NoState' responses, give up early
1816     # the daemon is alive but can't connect to the coordination server.
1817     if echo "$ts_output" | grep -q "NoState"; then
1818         consecutive_failures=$((consecutive_failures + 1))
1819         if [ "$consecutive_failures" -ge 40 ]; then
1820             echo ""
1821             echo "   [attempt $i/60] Stuck in 'NoState' for $consecutive_failures checks."
1822             echo "   Tailscale daemon is running but can't reach the coordination server."
1823             break
1824         fi
1825     else
1826         consecutive_failures=0
1827     fi
1828
1829     # Print a condensed status every 10 seconds so the user can see progress
1830     if [ $((i % 10)) -eq 0 ]; then
1831         ts_short=$(echo "$ts_output" | head -1 | tr -d '\r\n')
1832         echo "   [attempt $i/60] $ts_short"
1833     elif [ $((i % 5)) -eq 0 ]; then
1834         echo -n "$[i]"
1835     else
1836         echo -n "."
1837     fi
1838     sleep 1
1839 done
1840 echo ""
1841
1842 if [ "$connected" = "true" ]; then
1843     echo "Gateway container is online on the Tailscale mesh."
1844 elif [ "$BYPASS_PING" = "true" ]; then
1845     echo "[Warning] Gateway container check timed out. Proceeding anyway (GATEWAY_BYPASS_PING is true)..."
1846 else
1847     echo "ERROR: Gateway container failed to initialize Tailscale (timed out after 60s)." >&2
1848     echo " The container was seen in the following states during the wait:" >&2
1849     echo "$ts_output" | sed 's/~/ /' >&2

```

```

1850 echo "" >&2
1851 echo " This could mean:" >&2
1852 echo " 1. Tailscale auth key has expired check TS_AUTHKEY in .env" >&2
1853 echo " 2. Network issue inside the container check: docker compose logs tailscale" >&2
1854 echo " 3. gluetun VPN tunnel not connecting check: docker compose logs warp | tail -20" >&2
1855 echo "" >&2
1856 echo " Diagnose manually:" >&2
1857 echo " docker compose exec -T tailscale tailscale status" >&2
1858 echo " docker compose exec -T tailscale tailscale netcheck" >&2
1859 cleanup_handler ERR $LINENO
1860 exit 1
1861 fi
1862
1863 # 6b. Resolve the gateway's Tailscale IP.
1864 # Dynamically query the container for the IP (handles IP changes after re-auth)
1865 # and cache it in .gateway_ip for fast retrieval by recover.sh during Wi-Fi roaming.
1866 #
1867 # CRITICAL: `docker compose exec -T` on macOS/Colima injects carriage
1868 # returns (\r) into captured output. If $TS_IP ends with \r, then
1869 # `networksetup -setdnsservers` silently rejects the malformed IP and
1870 # DNS stays at 1.1.1.1 the primary bug this project was hitting.
1871 # `tr -d '\r'` strips the CR; `awk 'NR==1{print $1; exit}'` takes only
1872 # the first field of the first line.
1873 TS_IP=""
1874 if [ "$connected" = "true" ]; then
1875     TS_IP=$(run_with_timeout 5 docker compose exec -T tailscale tailscale ip -4 2>> output.log | tr -d '\r' |
1876         awk 'NR==1{print $1; exit}' || true)
1877     if [ -n "$TS_IP" ]; then
1878         echo "Resolved gateway Tailscale IP dynamically: $TS_IP"
1879         echo "$TS_IP" > .gateway_ip
1880     else
1881         echo " [!] Failed to resolve gateway IP dynamically. Trying fallback cache..." >&2
1882         TS_IP=$(read_adguard_ip || true)
1883     fi
1884 else
1885     TS_IP=$(read_adguard_ip || true)
1886 fi
1887
1888 # 7. Connect host Mac to Tailscale mesh.
1889 # With the standalone tailscaled daemon (registered as a per-user LaunchAgent or
1890 # system LaunchDaemon via 'brew services start tailscale'), the previous five-
1891 # phase flow collapses into two trivial CLI calls. There is no .app GUI to launch,
1892 # no menu bar item to click, and no Accessibility permission required.
1893 #
1894 # CRITICAL: If tailscaled is not running, `tailscale up` will silently fail.
1895 # We auto-start it before proceeding, and SKIP the rest of the Tailscale setup
1896 # if it can't be started (DNS hijack to a 100.x.x.x IP and exit-node routing
1897 # are impossible without the host on the mesh).
1898 #
1899 # We connect in two phases:
1900 # Phase A Join mesh WITHOUT exit node (so pre-flight checks can run)
1901 # Phase B Set exit node only after pre-flight checks confirm the gateway works
1902 HOST_ON_MESH=false
1903 SKIP_EXIT_NODE=false
1904 if [ "$HIJACK_HOST" = "false" ]; then
1905     echo -e "\n[Info] HIJACK_HOST is false (Headless Mode). Host networking will remain untouched."
1906     SKIP_EXIT_NODE=true
1907 elif [ -n "$TS_BIN" ]; then
1908     echo -n "Verifying tailscaled is reachable"
1909     daemon_ready=false
1910     status_exit=0
1911     for i in {1..10}; do
1912         # Test if the daemon responds to status check.
1913         # If status returns non-zero, check if it was a normal "stopped/disconnected" state
1914         # (exit code 1 is normal for disconnected). If the command exited quickly (not killed by timeout),
1915         # and the error is just "Tailscale is stopped", then the daemon is alive and ready.
1916         # But if it timed out (exit code 143) or is completely unresponsive, it's wedged.
1917         status_exit=0
1918         if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
1919             daemon_ready=true
1920             break
1921         else
1922             status_exit=$?
1923             if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
1924                 daemon_ready=true
1925                 break
1926             fi
1927         fi
1928     done
1929     echo -n "."
1930     sleep 1
1931 done

```

```

1930 echo ""
1931
1932 # If the daemon was unresponsive or socket check failed, attempt auto-restart
1933 if [ "$daemon_ready" != "true" ]; then
1934     if [ "$status_exit" -eq 143 ]; then
1935         warn "tailscaled daemon is wedged/unresponsive. Restarting it..."
1936     else
1937         warn "tailscaled daemon is not running. Attempting to start it..."
1938     fi
1939     restart_tailscaled_daemon
1940
1941     # Re-verify
1942     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1 || [ -S /var/run/tailscaled.socket ]; then
1943         daemon_ready=true
1944     fi
1945 fi
1946
1947 if [ "$daemon_ready" = "true" ]; then
1948     echo "tailscaled is responsive (running as a system service)."

```

```

2011     echo "PASS"
2012 else
2013     echo "FAIL"
2014     SKIP_EXIT_NODE=true
2015 fi
2016 elif command -v nslookup &>/dev/null; then
2017     if nslookup -port=${DNS_PROXY_PORT} google.com 127.0.0.1 &>/dev/null; then
2018         echo "PASS"
2019     else
2020         echo "FAIL"
2021         SKIP_EXIT_NODE=true
2022     fi
2023 else
2024     echo "SKIP (dig/nslookup not available)"
2025 fi
2026
2027 # Check 3: Does the WARP container have external internet?
2028 echo -n " [3/3] WARP container internet... "
2029 tmp_err=$(mktemp)
2030 if docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace >/dev/null 2>
2031     "$tmp_err"; then
2032     echo "PASS"
2033 else
2034     echo "FAIL"
2035     if [ -s "$tmp_err" ]; then
2036         err_msg=$(cat "$tmp_err" | tr -d '\r')
2037         echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Check 3] WARP container check failed: $err_msg" >> output.log
2038     fi
2039     SKIP_EXIT_NODE=true
2040 fi
2041 rm -f "$tmp_err"
2042
2043 # Act based on check results
2044 if [ "$SKIP_EXIT_NODE" != "true" ]; then
2045     echo -e "\nAll checks passed. Enabling exit node $TS_IP..."
2046     $TS_BIN up || true
2047     if $TS_BIN set --ssh=true --accept-dns=false --accept-routes=true --exit-node="$TS_IP" --exit-node-allow-
lan-access=true; then
2048         echo "Exit node enabled."
2049         add_warp_bypass_routes
2050         setup_exit_node_routing
2051         enable_killswitch
2052     else
2053         echo "[Warning] Failed to set exit node."
2054         SKIP_EXIT_NODE=true
2055     fi
2056 fi
2057
2058 if [ "$SKIP_EXIT_NODE" = "true" ]; then
2059     warn "Pre-flight checks failed EXIT NODE + DNS HIJACK SKIPPED."
2060     echo " Host stays on Tailscale mesh but your internet routes through your normal connection."
2061     echo " Troubleshoot with:"
2062     echo "     docker compose logs warp | tail -20"
2063     echo "     docker compose exec -T warp wget -qO- --timeout=5 https://www.cloudflare.com/cdn-cgi/trace"
2064     echo " Falling back to SOCKS5 proxy for traffic routing..."
2065 fi
2066 elif [ "$HOST_ON_MESH" = "true" ] && [ "$USE_EXIT_NODE" = "false" ]; then
2067     echo "USE_EXIT_NODE is false. Skipping exit node and DNS hijack."
2068     SKIP_EXIT_NODE=true
2069 fi
2070
2071 # 8. Apply DNS Hijacking.
2072 # If exit node is enabled: hijack DNS to gateway's Tailscale IP (routes through tailnet).
2073 # Otherwise, leave DNS at 1.1.1.1 (already set by reset at the top).
2074 # AdGuard is also available at localhost:${DNS_PROXY_PORT} for manual configuration in apps.
2075 if [ -n "$TS_IP" ] && [ "$HOST_ON_MESH" = "true" ] && [ "$SKIP_EXIT_NODE" != "true" ]; then
2076     echo -e "\nHijacking host DNS to point ONLY at AdGuard Home ($TS_IP)..."
2077     echo "Setting MagicDNS search domain 'ts.net' so tailnet hostnames resolve..."
2078     if force_dns_to_gateway "$TS_IP"; then
2079         echo "DNS hijack verified: single server = $TS_IP."
2080     else
2081         echo "[Warning] DNS hijack could not be verified."
2082         echo " If DNS is broken, run manually:"
2083         echo "     networksetup -setdnsservers \"\$ACTIVE_SERVICE\" \"$TS_IP"
2084         echo "     networksetup -setsearchdomains \"\$ACTIVE_SERVICE\" ts.net"
2085     fi
2086 elif [ -n "$TS_IP" ] && [ "$HIJACK_HOST" != "false" ]; then
2087     echo -e "\n[Info] Exit node not enabled (pre-flight checks failed)."
2088     echo " Trying local DNS proxy for ad-blocking..."
2089     if start_local_dns_proxy; then

```

```

2090     echo " Ad-blocking active via local DNS proxy through AdGuard."
2091 else
2092     echo " Local DNS proxy unavailable. DNS stays at 1.1.1.1."
2093     echo " AdGuard available at http://localhost:80 (port ${DNS_PROXY_PORT} for DNS via TCP)."
2094 fi
2095
2096 # Enable SOCKS5 proxy for traffic routing through WARP
2097 # (enabled whenever exit node is skipped, regardless of HOST_ON_MESH)
2098 echo -e "\n Enabling SOCKS5 proxy for TCP traffic routing through gateway WARP tunnel..."
2099 echo " (SOCKS5 proxy runs inside gateway container, routes through tun0 -> WARP)"
2100 enable_socks_proxy "$ACTIVE_SERVICE"
2101
2102 # Verify the SOCKS5 proxy works through WARP
2103 echo -n " Verifying SOCKS5 proxy routes through WARP... "
2104 if curl --socks5-hostname 127.0.0.1:${SOCKS_PROXY_PORT} --max-time 10 -s https://www.cloudflare.com/cdn-cgi/
2105     trace 2>/dev/null | grep -q "warp=on"; then
2106     echo "ok (warp=on)."
2107     echo " All TCP traffic now encrypted through Cloudflare WARP tunnel."
2108 else
2109     echo "FAILED (proxy not routing through WARP)."
2110     echo " Disabling SOCKS5 proxy internet stays on direct connection."
2111     disable_socks_proxy
2112     echo " SOCKS proxy available at localhost:${SOCKS_PROXY_PORT} for manual configuration."
2113 fi
2114 else
2115     echo -e "\n[Warning] Could not resolve gateway Tailscale IP."
2116 fi
2117
2118 # 9. Verify host connectivity
2119 if [ "$HOST_ON_MESH" = "true" ] && [ -n "$TS_BIN" ]; then
2120     if run_with_timeout 5 $TS_BIN status >> output.log 2>&1; then
2121         echo "Host is online on the Tailscale mesh."
2122     else
2123         echo "[Warning] Host is not responding on Tailscale."
2124     fi
2125 fi
2126
2127 ELAPSED=$(( SECONDS - TOGGLE_START_TIME ))
2128 echo -e "\nGateway has been successfully STARTED in ${ELAPSED} seconds."
2129
2130 # Prevent system from going to sleep while gateway is running
2131 start_sleep_prevention
2132
2133 if [[ "$HIJACK_HOST" != "false" ]]; then
2134     if [ -n "$TS_IP" ] && [ "$TS_IP" != "1.1.1.1" ]; then
2135         start_dns_watcher "$TS_IP"
2136     fi
2137
2138     # Start WARP liveness monitor (logs to output.log on warp flip)
2139     start_warp_watcher
2140 fi
2141
2142 # Reset sharing services after network configuration to prevent AirDrop freezes
2143 echo "Resetting macOS sharing services (AirDrop/AirPlay)..."
2144 reset_sharing_services
2145
2146 # Tell external watchers (post-wake / network-change) the gateway is up.
2147 write_gateway_active_marker
2148
2149 # Local Network Surveillance Check
2150 echo -e "\n"
2151 echo "Local Network Surveillance Check"
2152 echo ""
2153 # Count populated ARP cache entries to detect network scanning or high congestion
2154 # Using '-an' avoids hanging on DNS reverse resolution when the network state is in transition.
2155 ARP_COUNT=$(arp -an 2>/dev/null | grep -iv 'incomplete' | wc -l | tr -d ' ')
2156 if [ "$ARP_COUNT" -gt 15 ]; then
2157     warn "High ARP activity detected ($ARP_COUNT devices in cache)."
2158     warn "This Wi-Fi network may be heavily congested or actively scanned (e.g., arp-scan)."
2159     echo -e " [] Your traffic remains fully encrypted and invisible.\n"
2160 else
2161     echo -e " [] Local network looks quiet ($ARP_COUNT devices). No aggressive scanning detected.\n"
2162 fi
2163 fi
2164
2165 SUCCESS_RUN=true
2166
2167 # Final DNS state summary
2168 if [ -n "$ACTIVE_SERVICE" ]; then
2169     FINAL_DNS=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log || true)
2170     echo -e "\n"

```



```

2170     echo -e "DNS STATE: $FINAL_DNS"
2171     echo -e ""
2172 fi
2173
2174 if [ "$START_GATEWAY" = "true" ] && command -v docker >/dev/null 2>&1; then
2175     if docker compose logs rule-compiler 2>/dev/null | grep -q "Warning: Failed to fetch"; then
2176         echo -e "\n[Warning] One or more blacklist URLs failed to download (404/Offline)."
2177         echo "   The gateway gracefully fell back to yesterday's cached blacklist."
2178         echo "   (Run 'docker logs rule-compiler' later to investigate which link died.)"
2179     fi
2180 fi
2181
2182 rm -f "${LOCK_FILE:-/tmp/nullexit-toggle.lock}"
2183 if [ -t 0 ]; then
2184     read -rp "Press [r] and Enter to instantly reverse state, or just press Enter to exit: " USER_CHOICE
2185     if [[ "${USER_CHOICE}" == "r" || "${USER_CHOICE}" == "R" ]]; then
2186         echo "Reversing gateway state..."
2187         exec bash "$0"
2188     fi
2189 fi
2190
2191 echo -e "\nYou can close this terminal window now."

```

## 7 recover.sh

```

1 #!/bin/bash
2 # recover.sh nullexit KILL-SWITCH recovery script (macOS + Linux)
3 # Fixes internet when toggle.sh leaves you stranded.
4 #
5 # This is a NUCLEAR option: it undoes EVERYTHING toggle.sh does by:
6 #   1. Disconnecting host Tailscale from the mesh (exit-node off)
7 #   2. Resetting DNS to default on ALL network services
8 #   3. Disabling rogue proxy settings (SOCKS, web, secure web)
9 #   4. Flushing the DNS cache
10 #   5. Flushing stale routing table entries
11 #   6. Stopping gateway Docker containers
12 #   7. Power-cycling Wi-Fi
13 #   8. Verifying internet is working
14 #
15 # One file, both OSes: the dispatch below runs the self-contained Linux
16 # recovery and exits; on macOS it falls through to the richer dual-mode
17 # implementation (default teardown / --post-wake refresh).
18 #
19 # Usage:  double-click Recover-Gateway.applescript
20 #         OR  ./recover.sh
21 #         OR  ./recover.sh --post-wake    (macOS only  lighter refresh, keeps gateway live)
22
23 set -e
24
25 # Resolve script directory (used for path-relative refs)
26 # recover.sh lives at the repo root; SCRIPT_DIR lets us launch toggle.sh
27 # (and reference any other repo-root file) from the same directory no
28 # matter where the user invoked recover.sh from.
29 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
30 cd "$SCRIPT_DIR" || exit 1
31
32 #
33 # LINUX RECOVERY self-contained; runs the full teardown then exits before the
34 # macOS body below. Native Linux tooling (resolvectl / nmcli / ip) throughout.
35 #
36 if [[ "$OSTYPE" != "darwin"* ]]; then
37     # Source common formatting and helper functions
38     source "$SCRIPT_DIR/scripts/common.sh"
39
40     # Always add Homebrew path
41     export PATH="/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:$PATH"
42
43     echo ""
44     echo -e "${RED}${BOLD}${NC}"
45     echo -e "${RED}${BOLD}      Nullexit Gateway  RECOVERY MODE      ${NC}"
46     echo -e "${RED}${BOLD}${NC}"
47     echo ""
48     echo "This script will tear down the gateway and restore normal internet."
49     echo ""
50
51     # Cache sudo credentials upfront

```

```

52 # Prevents the script from hanging halfway through when it needs to run
53 # privileged commands (route flush, Wi-Fi power cycle, etc.).
54 echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
55 sudo -v
56 (
57     while true; do
58         sudo -n true 2>/dev/null
59         sleep 60
60     done
61 ) &
62 SUDO_KEEPER_PID=$!
63 trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
64
65 # 1. Disconnect host Tailscale
66 step "Disconnecting host Tailscale from mesh"
67 if command -v tailscale >> output.log 2>&1; then
68     # Check if actually connected first (with timeout tailscaled could be wedged)
69     if run_with_timeout 5 tailscale status >> output.log 2>&1; then
70         # Also explicitly reset any exit-node so it doesn't linger.
71         if run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= >> output.log
72             2>&1; then
73             ok "Exit-node preference cleared"
74         else
75             warn "tailscale up --reset didn't respond (tailscaled may be wedged)"
76         fi
77         ts_args=""
78         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
79         if run_with_timeout 10 tailscale down $ts_args >> output.log 2>&1; then
80             ok "Tailscale disconnected"
81         else
82             warn "tailscale down didn't respond"
83         fi
84     else
85         warn "tailscaled not reachable skipping (timeout after 5s)"
86     fi
87 else
88     warn "tailscale CLI not found skipping"
89 fi
90
91 # 2. Reset DNS to default on ALL network services
92 step "Resetting DNS to default on active interface"
93
94 ACTIVE_SERVICE=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
95 if [ -n "$ACTIVE_SERVICE" ]; then
96     sudo resolvectl revert "$ACTIVE_SERVICE" 2>> output.log || true
97     ok "DNS reset on $ACTIVE_SERVICE"
98 else
99     warn "Could not determine active network interface"
100 fi
101
102 # 3. Disable rogue proxy settings
103 step "Disabling proxy settings (SOCKS, web, secure web)"
104 echo " (Linux global SOCKS proxy configuration skipped.)"
105 ok "Proxy settings cleared"
106
107 # 4. Flush DNS cache
108 step "Flushing DNS cache"
109 if command -v resolvectl >> output.log 2>&1; then
110     sudo resolvectl flush-caches 2>> output.log || true
111     ok "DNS cache flushed"
112 else
113     warn "resolvectl not found"
114 fi
115
116 # 5. Clear stale routing table
117 step "Flushing stale routing table entries"
118 sudo ip route flush cache >> output.log 2>&1 || true
119 ok "Routing table flushed"
120
121 # 6. Stop gateway Docker containers
122 step "Stopping gateway Docker containers"
123 if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
124     if [ -f "docker-compose.yml" ]; then
125         docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
126     else
127         warn "docker-compose.yml not found in current directory"
128     fi
129 else
130     warn "Docker not available skipping"
131 fi

```

```

132
133 # 6b. Stopping sleep prevention (systemd-inhibit)
134 step "Stopping sleep prevention"
135 PID_FILE="/tmp/nullexit-caffeinate.pid"
136 if [ -f "$PID_FILE" ]; then
137     INHIBIT_PID=$(cat "$PID_FILE")
138     if [ -n "$INHIBIT_PID" ] && kill -0 "$INHIBIT_PID" 2>/dev/null; then
139         kill "$INHIBIT_PID" 2>/dev/null || true
140         ok "Sleep prevention stopped (PID $INHIBIT_PID)"
141     else
142         warn "Sleep prevention process not running, cleaning up stale PID file"
143     fi
144     rm -f "$PID_FILE"
145 else
146     ok "Sleep prevention was not active"
147 fi
148
149 # 7. Power-cycle Wi-Fi
150 step "Power-cycling Wi-Fi"
151 if command -v nmcli &>/dev/null; then
152     sudo nmcli radio wifi off 2>> output.log || true
153     sleep 2
154     sudo nmcli radio wifi on 2>> output.log || true
155     ok "Wi-Fi bounced via nmcli"
156     sleep 3
157 else
158     warn "nmcli not found. Falling back to ip link..."
159     WIFI_IFACE=$(ip link | grep -E '[0-9]+: (wl[a-zA-Z0-9]+|wlan[0-9]+):' | awk -F': ' '{print $2}')
160     if [ -n "$WIFI_IFACE" ]; then
161         sudo ip link set "$WIFI_IFACE" down 2>> output.log || true
162         sleep 2
163         sudo ip link set "$WIFI_IFACE" up 2>> output.log || true
164         ok "Wi-Fi bounced via ip link (interface $WIFI_IFACE)"
165         sleep 3
166     else
167         warn "No Wi-Fi interface detected."
168     fi
169 fi
170
171 # 8. Verify internet connectivity
172 verify_internet_connectivity
173
174 # Summary
175 echo ""
176 echo -e "${BOLD}${NC}"
177 echo -e "${BOLD}RECOVERY SUMMARY${NC}"
178 echo -e "${BOLD}${NC}"
179
180 # Show final DNS state
181 FINAL_DNS=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | tr
182 '\n' ' ' || echo "N/A")
183
184 echo " DNS ($ACTIVE_SERVICE): $FINAL_DNS"
185
186 echo ""
187
188 if [ "$INTERNET_OK" = "true" ]; then
189     echo -e " ${GREEN}${BOLD} Internet is working.${NC}"
190 else
191     echo -e " ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
192     echo " This might mean:"
193     echo " - Your Wi-Fi is disconnected from the router"
194     echo " - Tailscale is still interfering (try restarting tailscaled:"
195     echo "     systemctl restart tailscaled)"
196     echo " - You need to re-connect to Wi-Fi in the menu bar"
197     echo " - Try toggling Wi-Fi off/on in the menu bar"
198     echo ""
199     echo " Or try manually:"
200     echo "     resolvectl revert $ACTIVE_SERVICE"
201     extra_args=""
202     if is_kill_switch_enabled; then extra_args="--accept-risk=lose-ssh"; fi
203     echo "     tailscale down $extra_args"
204     echo "     sudo ip route flush cache"
205     echo "     systemctl restart tailscaled"
206 fi
207
208 echo ""
209 echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
210 echo ""
211 exit 0
212 fi

```

```

212 #
213 #
214 # macOS RECOVERY (dual-mode: default teardown / --post-wake refresh)
215 #
216 # Enforce Cryptographic Script Integrity
217 if [ -f "scripts/crypto.sh" ]; then
218     if ! bash scripts/crypto.sh --verify; then
219         exit 1
220     fi
221 fi
222
223 # Argument parsing
224 # --post-wake: lightweight refresh after sleep/wake or Wi-Fi roam.
225 # Keeps the gateway live while HUPing DNS caches, refreshing the
226 # Tailscale exit-node leak, re-hijacking host DNS, and force-
227 # recreating the warp container if its gluetun healthcheck failed.
228 # Does NOT tear down Docker, Tailscale, sleep prevention, or Wi-Fi.
229 # Called from scripts/watcher.sh on every wake / network-change event
230 # while /tmp/nullexit-gateway-active.marker is present (i.e. the
231 # gateway is currently up). See devref.md 10.29 for the why.
232 POST_WAKE=false
233 AUTO_MODE=false
234 for arg in "$@"; do
235     if [ "$arg" = "--post-wake" ]; then
236         POST_WAKE=true
237     elif [ "$arg" = "--auto" ] || [ "$arg" = "--non-interactive" ]; then
238         AUTO_MODE=true
239     fi
240 done
241
242 rm -f "$SCRIPT_DIR/TUNNEL_FAILED_CLOSED.marker"
243
244 # Prevent concurrent execution of lifecycle scripts (toggle.sh / recover.sh)
245 LOCK_FILE="/tmp/nullexit-toggle.lock"
246 if [ -f "$LOCK_FILE" ]; then
247     LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
248     if [ -n "$LOCK_PID" ] && [ "$LOCK_PID" != "$$" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
249         if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -q -E "toggle\.sh|recover\.sh"; then
250             if [ "$POST_WAKE" = "true" ]; then
251                 echo "[$(date -u +%FT%TZ)] POST-WAKE skipped: another lifecycle script (PID $LOCK_PID) is running." >>
252                     output.log
253                 exit 0
254             else
255                 echo "Error: Another instance of toggle.sh or recover.sh (PID $LOCK_PID) is already running." >&2
256                 exit 1
257             fi
258         fi
259     fi
260 fi
261 echo "$$" > "$LOCK_FILE"
262
263 # Clear the active-state marker in nuclear recovery mode since the gateway is being
264 # completely torn down. This prevents scripts/watcher.sh from triggering post-wake
265 # recoveries in response to network changes caused by the teardown itself.
266 if [ "$POST_WAKE" = "false" ]; then
267     rm -f "/tmp/nullexit-gateway-active.marker"
268 fi
269
270 # WARP Watcher inhibit marker: written immediately in --post-wake mode so the
271 # background WARP Watcher (in toggle.sh) skips failure counting while we are
272 # intentionally tearing down / recreating the warp container. Without this, the
273 # watcher accumulates 6 consecutive warp=off readings during force-recreate (~35s)
274 # and fires nuclear recover.sh, killing a gateway that would have been healthy.
275 WARP_INHIBIT_FILE="/tmp/nullexit-warp-inhibit.marker"
276 if [ "$POST_WAKE" = "true" ]; then
277     touch "$WARP_INHIBIT_FILE"
278     echo "[$(date -u +%FT%TZ)] POST-WAKE started WARP Watcher inhibited to prevent spurious shutdown during warp
279         recreate." >> output.log
280 fi
281
282 # Ensure the inhibit marker and lock file are always removed when the script exits (success or error)
283 cleanup_recover() {
284     rm -f "$WARP_INHIBIT_FILE"
285     if [ -f "$LOCK_FILE" ] && [ "$(cat "$LOCK_FILE" 2>/dev/null)" = "$$" ]; then
286         rm -f "$LOCK_FILE"
287     fi
288 }
289 trap cleanup_recover EXIT
290
291 # Source common formatting and helper functions

```

```

291 source "$SCRIPT_DIR/scripts/common.sh"
292
293 # Always add Homebrew path
294 setup_standard_path
295
296 echo ""
297 if [ "$POST_WAKE" = "true" ]; then
298     echo -e "${YELLOW}${BOLD}${NC}"
299     echo -e "${YELLOW}${BOLD} Nullexit POST-WAKE / POST-ROAM LIGHT ${NC}"
300     echo -e "${YELLOW}${BOLD} (keeps the gateway live, refreshes) ${NC}"
301     echo -e "${YELLOW}${BOLD}${NC}"
302     echo ""
303     echo "Re-hijacking DNS, refreshing Tailscale exit-node, force-recreating"
304     echo "warp if unhealthy, and resetting sharing services. The gateway stays"
305     echo "up. docker compose / Wi-Fi / caffeine are NOT touched."
306     echo ""
307 else
308     echo -e "${RED}${BOLD}${NC}"
309     echo -e "${RED}${BOLD} Nullexit Gateway RECOVERY MODE ${NC}"
310     echo -e "${RED}${BOLD}${NC}"
311     echo ""
312     echo "This script will tear down the gateway and restore normal internet."
313     echo ""
314 fi
315
316 # Cache sudo credentials upfront (default mode ONLY)
317 # Skip in --post-wake: the LaunchAgent-launched watcher has no TTY, so `sudo -v`
318 # would block forever waiting for a password. All post-wake commands use `sudo -n`
319 # which is non-blocking and relies on /etc/sudoers.d/nullexit NOPASSWD entries.
320 SUDO_KEEPER_PID=""
321 if [ "$POST_WAKE" = "false" ] && [ "$AUTO_MODE" = "false" ] && [ -t 0 ]; then
322     echo -e "${YELLOW}${BOLD}Authentication required for network recovery...${NC}"
323     sudo -v
324     (
325         while true; do
326             sudo -n true 2>/dev/null
327             sleep 60
328         done
329     ) &
330     SUDO_KEEPER_PID=$!
331     trap 'kill $SUDO_KEEPER_PID 2>/dev/null' EXIT
332 fi
333
334 # Resolve physical default interface and gateway upfront
335 PHYSICAL_IFACE=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
    getline; print $2; exit}')
336 [ -z "$PHYSICAL_IFACE" ] && PHYSICAL_IFACE="en0"
337
338 if [ "$POST_WAKE" = "true" ]; then
339     wait_for_dhcp_settle
340 else
341     # Quick resolution in non-post-wake mode (non-blocking)
342     PHYSICAL_GW=$(ipconfig getpacket "$PHYSICAL_IFACE" 2>> output.log | awk -F'[]{}' '/router/{print $2}')
343 fi
344
345 # 1. Disconnect host Tailscale (default) / Refresh exit-node (post-wake)
346
347 if [ "$POST_WAKE" = "true" ]; then
348     step "Refreshing host Tailscale exit-node routing (post-wake)"
349     if command -v tailscale >> output.log 2>&1; then
350         # Re-issue the SAME exit-node up command toggle.sh START uses. This refreshes
351         # the DERP relay mapping and re-asserts the exit-node preference without
352         # dropping the host's mesh connection (which `tailscale down` would).
353         TS_IP=""
354         for src in .gateway_ip; do
355             [ -f "$src" ] || continue
356             TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
357             [ -n "$TS_IP" ] && break
358         done
359
360         if [ -n "$TS_IP" ]; then
361             step "Re-establishing exit node $TS_IP via 'tailscale set --exit-node'"
362             # `tailscale set` only mutates the named preference (exit-node here).
363             # Crucially it does NOT call --reset, which would re-apply ALL flags and
364             # trigger a sub-second route-table re-evaluation cycle each post-wake.
365             # On already-connected nodes, the lighter `set` keeps the existing tunnel
366             # alive while only changing the exit-node preference.
367             if run_with_timeout 15 tailscale set --exit-node="$TS_IP" \
368                 --exit-node-allow-lan-access=true \
369                 >> output.log 2>&1; then
370                 ok "Exit-node $TS_IP re-asserted via 'tailscale set'"

```

```

371     else
372         warn "tailscale set --exit-node=$TS_IP didn't respond; falling back to mesh-only"
373         run_with_timeout 10 tailscale set --exit-node= \
374             >> output.log 2>&1 || true
375     fi
376 else
377     warn ".gateway_ip missing cannot re-assert exit node"
378     run_with_timeout 10 tailscale up --reset --ssh=true --accept-dns=false --exit-node= \
379         >> output.log 2>&1 || true
380 fi
381 else
382     warn "tailscale CLI not found"
383 fi
384 else
385     step "Disconnecting host Tailscale from mesh"
386     disconnect_tailscale_host "tailscale"
387 fi
388
389 # 2. Reset DNS to default on ALL network services (default) / Re-hijack gateway DNS (post-wake)
390 if [ "$POST_WAKE" = "true" ]; then
391     step "Re-hijacking host DNS to gateway IP (post-wake / post-roam)"
392     TS_IP=""
393     for src in .gateway_ip; do
394         [ -f "$src" ] || continue
395         TS_IP=$(cat "$src" 2>> output.log | tr -d '\r' | awk 'NR==1{print $1; exit}' || true)
396         [ -n "$TS_IP" ] && break
397     done
398
399     if [ -n "$TS_IP" ]; then
400         # Detect the active network service (using helper in common.sh)
401         ACTIVE_SVC=$(get_active_service)
402         ENO_SVC=$(get_eno_service)
403
404         networksetup -setsearchdomains "$ACTIVE_SVC" "ts.net" 2>> output.log || true
405         networksetup -setdnsservers "$ACTIVE_SVC" "$TS_IP" 2>> output.log || true
406         networksetup -setsearchdomains "$ENO_SVC" "ts.net" 2>> output.log || true
407         networksetup -setdnsservers "$ENO_SVC" "$TS_IP" 2>> output.log || true
408
409         HITS=$(networksetup -getdnsservers "$ACTIVE_SVC" 2>> output.log | grep -Fx "$TS_IP" || true)
410         if [ -n "$HITS" ]; then
411             ok "DNS re-hijacked to $TS_IP on services: $ACTIVE_SVC, $ENO_SVC"
412         else
413             warn "networksetup accepted but DNS read-back didn't include $TS_IP"
414         fi
415     else
416         warn ".gateway_ip missing cannot re-hijack DNS (user must run toggle.sh START)"
417     fi
418 else
419     step "Resetting DNS to default on all network services (empty = DHCP)"
420
421     # Get ALL network services (handles spaces in names, skips disabled ones)
422     SERVICES=$(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
423         true)
424
425     if [ -z "$SERVICES" ]; then
426         # Fallback to common names
427         SERVICES="Wi-Fi Ethernet Thunderbolt Ethernet USB 10/100 LAN"
428     fi
429
430     RESET_COUNT=0
431     while IFS= read -r service; do
432         # Strip leading asterisk (disabled services)
433         service_clean=$(echo "$service" | sed 's/^\*//')
434         [ -z "$service_clean" ] && continue
435
436         # Check if this service has a hardware port (skip VPN/service-only entries)
437         if networksetup -listallhardwareports 2>> output.log | grep -A1 "Port: $service_clean$" | grep -q "Device:"
438         || [ "$service_clean" = "Wi-Fi" ]; then
439             (
440                 networksetup -setsearchdomains "$service_clean" "Empty" 2>> output.log
441                 # Set to "empty" so macOS uses DHCP-assigned DNS (not hardcoded 1.1.1.1)
442                 sudo networksetup -setdnsservers "$service_clean" "empty" 2>> output.log
443             ) && RESET_COUNT=$((RESET_COUNT + 1)) || true
444         fi
445     done <<< "$SERVICES"
446
447     ok "DNS reset on $RESET_COUNT service(s)"
448 fi
449
450 # 3. Disable rogue proxy settings (default only gateway uses proxy in non-exit-node path)
451 if [ "$POST_WAKE" = "false" ]; then

```



```

450 step "Disabling proxy settings (SOCKS, web, secure web)"
451 for svc in $(networksetup -listallnetworkservices 2>> output.log | grep -v "^An asterisk" | grep -v "^$" ||
452     echo "Wi-Fi"); do
453     svc_clean=$(echo "$svc" | sed 's/^\\*/\\')
454     [ -z "$svc_clean" ] && continue
455     disable_all_proxies "$svc_clean"
456 done
457 ok "Proxy settings cleared"
458 else
459 step "Proxy settings: leaving untouched (post-wake keeps gateway SOCKS wiring alive)"
460 fi
461 # 4. Flush macOS DNS cache (always safe and useful in both modes)
462 step "Flushing DNS cache"
463 if command -v dscacheutil >> output.log 2>&1; then
464     flush_dns_cache
465     ok "DNS cache flushed (mDNSResponder HUP)"
466 else
467     warn "dscacheutil not found"
468 fi
469 # 5. Clear stale routing table (always safe in both modes)
470 step "Flushing stale routing table entries"
471 # Resolve physical default gateway IP before flushing
472 # We cannot trust `route get default` here because Tailscale may have already hijacked
473 # the default route to utunX. We must read the actual DHCP router assignment from the hardware.
474 # PHYSICAL_IFACE and PHYSICAL_GW resolved upfront
475 # Instead of full route -n flush (which destroys macOS loopback/multicast and wedges the OS until reboot),
476 # we cleanly delete the Tailscale overrides and Cloudflare WARP bypass routes in ALL modes.
477 # This ensures tailscaled isn't blocked from reaching its control plane when waking from sleep.
478 sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true
479 sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
480 remove_warp_bypass_routes
481 ok "Routing overrides cleared"
482
483 if [ "$POST_WAKE" = "false" ]; then
484     disable_killswitch
485     ok "Routing table flush completed via targeted deletes and firewall disabled"
486 else
487     ok "Routing overrides cleared (post-wake mode to preserve Colima bridge100)"
488 fi
489
490 # In post-wake mode we don't bounce Wi-Fi, so we MUST manually restore the physical default route
491 if [ "$POST_WAKE" = "true" ] && [ -n "$PHYSICAL_GW" ]; then
492     # Ensure physical gateway is set just in case it was dropped
493     sudo -n route add default "$PHYSICAL_GW" >> output.log 2>&1 || true
494 fi
495
496 # CONDITIONAL BYPASS ROUTE RESTORATION:
497 # If this was a post-wake recovery and the exit node is active, the route flush
498 # just wiped the static bypass routes for the WARP endpoints. We must re-add
499 # them immediately to prevent a routing loop when Tailscale starts sending traffic.
500 if [ "$POST_WAKE" = "true" ]; then
501     if [ -z "${ACTIVE_IF:-}" ]; then
502         ACTIVE_IF=$(route get default 2>> output.log | awk '/interface:/{print $2; exit}')
503         if [ -z "$ACTIVE_IF" ] || [[ "$ACTIVE_IF" =~ ^utun ]] || [[ "$ACTIVE_IF" =~ ^tun ]]; then
504             for i in en0 en1 en2 en3; do
505                 if ifconfig "$i" 2>> output.log | grep -q 'status: active'; then
506                     ACTIVE_IF="$i"; break
507                 fi
508             done
509         fi
510         [ -z "$ACTIVE_IF" ] && ACTIVE_IF="en0"
511     fi
512     echo "Re-adding host bypass routes for Cloudflare WARP endpoints..."
513     remove_warp_bypass_routes
514     if [ -n "${PHYSICAL_GW:-}" ]; then
515         add_warp_bypass_routes "$PHYSICAL_GW"
516     else
517         add_warp_bypass_routes "$ACTIVE_IF"
518     fi
519
520 # Also re-apply the default route pointing to the Tailscale interface
521 ts_iface=$(ifconfig 2>> output.log | grep -B4 "inet 100." | grep -E '[a-z0-9]+' | cut -d: -f1 | head -n 1)
522 if [ -n "$ts_iface" ]; then
523     echo "Re-routing default gateway to $ts_iface using 0.0.0.0/1 split..."
524     # DO NOT delete the physical default route! It crashes Colima.
525     # Use the /1 trick to mathematically override the default route.
526     sudo -n route delete -net 0.0.0.0/1 >> output.log 2>&1 || true

```

```

530 sudo -n route delete -net 128.0.0.0/1 >> output.log 2>&1 || true
531 sudo -n route add -net 0.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
532 sudo -n route add -net 128.0.0.0/1 -interface "$ts_iface" >> output.log 2>&1 || true
533 enable_killswitch
534 fi
535 fi
536
537 # 6. Stop gateway Docker containers (default) / Force-recreate warp if unhealthy (post-wake)
538 if [ "$POST_WAKE" = "true" ]; then
539 # 6a. Check for Roaming Network Mismatch (Bridged vs Enterprise Wi-Fi)
540 if [ -f "/tmp/nullexit_colima_mode.txt" ] && [ -f "$SCRIPT_DIR/../lan_p2p_detected" ]; then
541 current_colima_mode=$(cat /tmp/nullexit_colima_mode.txt)
542 p2p_allowed=$(cat "$SCRIPT_DIR/../lan_p2p_detected")
543 required_colima_mode="bridged"
544 [ "$p2p_allowed" == "false" ] && required_colima_mode="shared"
545
546 if [ "$current_colima_mode" != "$required_colima_mode" ]; then
547 warn "Network roaming mismatch detected!"
548 warn "Colima is running in '$current_colima_mode' mode but new network requires '$required_colima_mode'."
549 warn "Executing full gateway restart to protect against MAC-spoofing deauthentication..."
550 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Roam] Mismatch ($current_colima_mode vs $required_colima_mode).
Triggering toggle.sh --restart." >> output.log
551 nohup bash "$SCRIPT_DIR/../toggle.sh" --restart >/dev/null 2>&1 &
552 exit 0
553 fi
554 fi
555
556 step "Checking Colima VM state"
557 colima_status=$(colima status 2>&1)
558 if echo "$colima_status" | grep -qi "running"; then
559 ok "Colima VM is running"
560 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima status check passed." >> output.log
561 else
562 warn "Colima VM is NOT running or unresponsive"
563 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Colima is unhealthy or dead! Output: $colima_status"
>> output.log
564 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Attempting to restart Colima..." >> output.log
565 colima restart >> output.log 2>&1 || warn "Failed to restart Colima"
566 fi
567
568 step "Checking warp container health (gluetun UDP tunnel)"
569 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Inspecting warp container state..." >> output.log
570 # The warp container is the most failure-prone link in the chain at wake/roam:
571 # its UDP NAT binding to Cloudflare's edge (162.159.192.1:2408) silently dies
572 # during a Wi-Fi roam. We verify the tunnel is actually passing traffic via curl.
573 WARP_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
574 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container status: $WARP_STATE" >> output.log
575
576 if [ "$WARP_STATE" = "missing" ]; then
577 warn "warp container is missing leaving gateway containers alone (full relaunch requires \"$SCRIPT_DIR/
toggle.sh\")"
578 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container missing from Docker engine. Requires
full toggle.sh launch." >> output.log
579 else
580 tmp_err=$(mktemp)
581 if docker compose exec -T warp wget -qO- --timeout=3 https://www.cloudflare.com/cdn-cgi/trace 2>"$tmp_err"
| grep -q "warp=on"; then
582 ok "warp container is healthy and actively tunneling traffic (no recreate needed)"
583 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container passed live traffic healthcheck (
Cloudflare trace successful)." >> output.log
584 else
585 warn "warp container tunnel is dead (Docker status: $WARP_STATE) force-recreating all gateway containers
to restore network namespace"
586 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container is DEAD or STUCK. Forcing recreation
of gateway stack..." >> output.log
587 if [ -s "$tmp_err" ]; then
588 err_msg=$(cat "$tmp_err" | tr -d '\r')
589 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] Healthcheck failure details: $err_msg" >> output.
log
590 fi
591 docker compose up -d --force-recreate >> output.log 2>&1 || warn "force-recreate failed"
592
593 NEW_STATE=$(docker inspect --format '{{.State.Health.Status}}' warp 2>> output.log || echo "missing")
594 ok "warp container health after recreate: $NEW_STATE"
595 echo "[$(date '+%Y-%m-%d %H:%M:%S')] [Docker/Colima] warp container health after recreation: $NEW_STATE"
>> output.log
596 fi
597 rm -f "$tmp_err"
598 fi
599 else
600 step "Stopping gateway Docker containers"

```

```

601 if command -v docker >> output.log 2>&1 && docker info >> output.log 2>&1; then
602     if [ -f "docker-compose.yml" ]; then
603         docker compose down -t 5 2>> output.log && ok "Containers stopped" || warn "No containers were running"
604     else
605         warn "docker-compose.yml not found in current directory"
606     fi
607 else
608     warn "Docker not available skipping"
609 fi
610 fi
611
612 # 6b. Stopping sleep prevention (caffeinate) default only
613 if [ "$POST_WAKE" = "false" ]; then
614     step "Stopping sleep prevention"
615     stop_pidfile_daemon "$PID_CAFFEINATE" "system sleep prevention"
616 else
617     step "Sleep prevention: leaving untouched (caffeinate already re-acquired by parent shell)"
618 fi
619
620 # 6c. Stopping DNS Watcher default only
621 if [ "$POST_WAKE" = "false" ]; then
622     step "Stopping DNS Watcher"
623     stop_pidfile_daemon "$PID_DNS_WATCHER" "background DNS Watcher"
624 else
625     step "DNS Watcher: leaving untouched (it keeps re-hijacking DNS on every roam)"
626 fi
627
628 # 7. Power-cycle Wi-Fi default only
629 if [ "$POST_WAKE" = "false" ]; then
630     step "Power-cycling Wi-Fi"
631     bounce_wifi_interfaces
632 else
633     step "Wi-Fi: leaving untouched (post-wake keeps the current Wi-Fi association)"
634 fi
635
636 # 8. Verify (default checks direct internet; post-wake verifies the gateway)
637 if [ "$POST_WAKE" = "true" ]; then
638     step "Verifying gateway is healthy after post-wake refresh"
639
640     # Wait a moment for tailscale + warp healthcheck to settle
641     sleep 3
642
643     GATEWAY_OK=false
644     if command -v dig >> output.log 2>&1; then
645         # Source procedural ports if available
646         if [ -f "/tmp/nullexit-ports.env" ]; then
647             source "/tmp/nullexit-ports.env"
648         fi
649         DNS_PORT=${DNS_PROXY_PORT:-5354}
650
651         # AdGuard is exposed at localhost:${DNS_PORT}. If this resolves, the warp tunnel
652         # is functioning properly because AdGuard routes upstream to warp.
653         if dig +tcp @127.0.0.1 -p "$DNS_PORT" google.com +short +timeout=5 >> output.log 2>&1; then
654             ok "Gateway DNS (AdGuard via localhost:${DNS_PORT}) works"
655             GATEWAY_OK=true
656         else
657             warn "Gateway DNS not responding yet"
658         fi
659     fi
660
661     # Direct host gateway check via Tailscale ping (relayed pong counts as success)
662     if command -v tailscale >> output.log 2>&1; then
663         TS_IP=""
664         [ -f .gateway_ip ] && TS_IP=$(read_adguard_ip || true)
665         if [ -n "$TS_IP" ] && tailscale ping --until-direct=false -c 1 --timeout 4s "$TS_IP" 2>> output.log | grep
666             -q pong; then
667             ok "Gateway $TS_IP reachable via Tailscale"
668             GATEWAY_OK=true
669         else
670             warn "Tailscale ping to gateway did not pong exit-node may still be re-establishing"
671         fi
672     fi
673
674     if [ "$GATEWAY_OK" = "true" ]; then
675         echo -e " ${GREEN}${BOLD} Gateway is healthy after post-wake refresh.${NC}"
676     else
677         echo -e " ${YELLOW}${BOLD} Gateway recovery is in-progress.${NC}"
678         echo " The route may need another 30s before queries succeed."
679         echo " If it stays broken for >60s, run: bash \"\$SCRIPT_DIR/toggle.sh\""
680     fi
681 else

```

```

681     verify_internet_connectivity
682 fi
683
684 # Summary
685 echo ""
686 echo -e "${BOLD}${NC}"
687 if [ "$POST_WAKE" = "true" ]; then
688     echo -e "${BOLD}POST-WAKE SUMMARY${NC}"
689 else
690     echo -e "${BOLD}RECOVERY SUMMARY${NC}"
691 fi
692 echo -e "${BOLD}${NC}"
693
694 # Show final DNS state
695 FINAL_DNS=$(networksetup -getdnsservers "Wi-Fi" 2>/dev/null || echo "N/A")
696 echo "  DNS (Wi-Fi):  $FINAL_DNS"
697
698 FINAL_DNS2=$(networksetup -getdnsservers "Ethernet" 2>/dev/null || true)
699 if [[ -z "$FINAL_DNS2" || "$FINAL_DNS2" == *"not a recognized"* || "$FINAL_DNS2" == *"Error"* ]]; then
700     FINAL_DNS2="N/A (Not configured)"
701 fi
702 echo "  DNS (Ethernet): $FINAL_DNS2"
703
704 echo ""
705
706 if [ "$POST_WAKE" = "false" ]; then
707     if [ "$INTERNET_OK" = "true" ]; then
708         echo -e "  ${GREEN}${BOLD} Internet is working.${NC}"
709     else
710         echo -e "  ${YELLOW}${BOLD} Could not verify internet connectivity.${NC}"
711         echo "    This might mean:"
712         echo "    - Your Wi-Fi is disconnected from the router"
713         echo "    - Tailscale is still interfering (try restarting tailscaled:"
714         echo "      brew services restart tailscale)"
715         echo "    - You need to re-connect to Wi-Fi in the menu bar"
716         echo "    - Try toggling Wi-Fi off/on in the menu bar"
717         echo ""
718         echo "    Or try manually:"
719         echo "      networksetup -setdnsservers Wi-Fi empty"
720         echo "      tailscale down"
721         echo "      sudo route delete -net 0.0.0.0/1"
722         echo "      sudo route delete -net 128.0.0.0/1"
723         echo "      brew services restart tailscale"
724     fi
725 fi
726
727 # Reset sharing services to prevent AirDrop freezes after interface changes
728 # This is the SAME step in BOTH modes (post-wake inherits the previous fix from 10.27).
729 echo -e "  Resetting macOS sharing services (AirDrop/AirPlay)..."
730 reset_sharing_services
731
732 echo ""
733 if [ "$POST_WAKE" = "true" ]; then
734     echo -e "${YELLOW}${BOLD}Post-wake refresh complete.${NC}"
735     # Remove the WARP Watcher inhibit marker now that the gateway is fully healthy.
736     # The watcher will resume normal liveness monitoring on its next 5s poll.
737     rm -f "$WARP_INHIBIT_FILE"
738     echo "[$(date -u +%FT%TZ)] POST-WAKE complete  WARP Watcher inhibit released." >> output.log
739 else
740     echo -e "${GREEN}${BOLD}Recovery complete.${NC}"
741 fi
742 echo ""

```

## 8 setup.sh

```

1  #!/bin/bash
2  # setup.sh nullexit one-shot setup script (macOS)
3  # Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
4  set -euo pipefail
5
6  # Source common formatting and helper functions
7  source "$(dirname "$0")/scripts/common.sh"
8
9  # Run common installation logic
10 source "$(dirname "$0")/scripts/setup-common.sh"
11

```

```

12 # Compile macOS Shortcuts
13 echo -e "\n${BOLD}Compiling macOS Desktop Shortcuts...${NC}"
14 if command -v osacompile &> /dev/null; then
15     osacompile -o "Toggle Gateway.app" "Toggle-Gateway.applescript" >/dev/null 2>&1
16     osacompile -o "Recover Gateway.app" "Recover-Gateway.applescript" >/dev/null 2>&1
17     echo "    Created 'Toggle Gateway.app' and 'Recover Gateway.app'"
18 fi
19
20 # Done
21 echo ""
22 echo -e "${GREEN}${BOLD}${NC}"
23 echo -e "${GREEN}${BOLD}                Setup complete!                ${NC}"
24 echo -e "${GREEN}${BOLD}${NC}"
25 echo ""
26 echo -e "${BOLD}One manual step remaining approve the exit node:${NC}"
27 echo "    https://login.tailscale.com/admin/machines"
28 echo "    Find your gateway '...' 'Edit route settings' enable Exit Node"
29 echo ""
30
31 echo -e "${YELLOW}${BOLD}Sudoers / Background Execution Requirements:${NC}"
32 echo "    To allow silent, non-interactive execution of the firewall and network routes (especially during sleep/
wake recovery), you must configure passwordless sudo."
33 echo "    Run this exact command in your terminal (grants are scoped to least privilege):"
34 echo "        echo \"\$USER ALL=(root) NOPASSWD: /sbin/pfctl, /usr/sbin/networksetup, /sbin/route, /sbin/ifconfig,
/usr/bin/dscacheutil -flushcache, /usr/bin/killall -HUP mDNSResponder, /usr/bin/killall sharingd rapportd,
/usr/bin/pkill -9 tailscaled, /usr/bin/pkill -f dns-proxy.py, /bin/kill, /usr/bin/true, /opt/homebrew/bin/
brew services * tailscale, /opt/homebrew/bin/brew uninstall tailscale, /usr/local/bin/brew services *
tailscale, /usr/local/bin/brew uninstall tailscale, /usr/bin/python3 -I \$PWD/scripts/dns-proxy.py, /opt/
homebrew/bin/python3 -I \$PWD/scripts/dns-proxy.py, /usr/local/bin/python3 -I \$PWD/scripts/dns-proxy.py\"
| sudo tee /etc/sudoers.d/nullexit >/dev/null && sudo visudo -cf /etc/sudoers.d/nullexit"
35 echo ""
36
37 if [[ -n "${TS_IP:-}" ]]; then
38     echo -e "${BOLD}AdGuard Home dashboard:${NC}    http://${TS_IP}:3000    (username: admin)"
39     echo ""
40 fi
41
42 echo -e "${BOLD}To toggle the gateway on/off:${NC}"
43 echo "    macOS: double-click the 'Toggle Gateway' app icon!"
44 echo "    Windows: double-click Toggle-Gateway.bat"
45 echo ""
46
47 if [[ "$OSTYPE" == "darwin"* ]]; then
48     # LuLu localhost filtering check
49     if [ -d "/Applications/LuLu.app" ] || systemextensionsctl list 2>/dev/null | grep -q "com.objective-see.
lulu"; then
50         echo -e "${YELLOW}${BOLD}LuLu Firewall Detected:${NC}"
51         echo "    LuLu's localhost (loopback) filtering is OFF by default. To properly protect the"
52         echo "    local nullexit proxy ports from malicious local processes, you must manually enable"
53         echo "    'Allow Loopback' (or block it selectively) in LuLu's Preferences -> Rules."
54         echo ""
55     fi
56
57     echo -e "${YELLOW}${BOLD}Note on macOS Tailscale Sandboxing:${NC}"
58     echo "    The Tailscale GUI app (tailscale-app) installed on macOS is sandboxed."
59     echo "    While you cannot run a Tailscale SSH server on this Mac host, you can"
60     echo "    still SSH outbound to other devices on the mesh using their MagicDNS"
61     echo "    addresses directly (e.g. ssh user@device.tailnet-name.ts.net)."
62     echo ""
63 fi

```

## 9 scripts/setup-linux.sh

```

1 #!/bin/bash
2 # scripts/setup-linux.sh nullexit one-shot setup script (Linux)
3 # Configures Cloudflare WARP + Tailscale + AdGuard Home from scratch.
4 set -euo pipefail
5
6 # Source common formatting and helper functions
7 source "$(dirname "$0")/common.sh"
8
9 # Run common installation logic
10 source "$(dirname "$0")/setup-common.sh"
11
12 # Compile Linux Desktop Shortcuts
13 echo -e "\n${BOLD}Creating Linux Desktop Shortcuts...${NC}"

```

```

14 DIR="$(pwd)"
15 cat <<EOF > "Toggle Gateway.desktop"
16 [Desktop Entry]
17 Version=1.0
18 Name=Toggle Gateway
19 Comment=Start or Stop Nullexit Gateway
20 Exec=bash -c "cd '$DIR' && ./toggle.sh; read -p 'Press Enter to close...' dummy"
21 Icon=utilities-terminal
22 Terminal=true
23 Type=Application
24 Categories=Network;Security;
25 EOF
26
27 cat <<EOF > "Recover Gateway.desktop"
28 [Desktop Entry]
29 Version=1.0
30 Name=Recover Gateway
31 Comment=Emergency DNS and Network Recovery
32 Exec=bash -c "cd '$DIR' && ./recover.sh; read -p 'Press Enter to close...' dummy"
33 Icon=utilities-terminal
34 Terminal=true
35 Type=Application
36 Categories=Network;Security;
37 EOF
38 echo "    Created 'Toggle Gateway.desktop' and 'Recover Gateway.desktop'"
39
40 # Done
41 echo ""
42 echo -e "${GREEN}${BOLD}${NC}"
43 echo -e "${GREEN}${BOLD}                Setup complete!                ${NC}"
44 echo -e "${GREEN}${BOLD}${NC}"
45 echo ""
46 echo -e "${BOLD}One manual step remaining  approve the exit node:${NC}"
47 echo "    https://login.tailscale.com/admin/machines"
48 echo "    Find your gateway '...' 'Edit route settings'  enable Exit Node"
49 echo ""
50
51 if [[ -n "${TS_IP:-}" ]]; then
52     echo -e "${BOLD}AdGuard Home dashboard:${NC}  http://${TS_IP}:3000  (username: admin)"
53     echo ""
54 fi
55
56 echo -e "${BOLD}To toggle the gateway on/off:${NC}"
57 echo "    Linux: double-click the 'Toggle Gateway' desktop icon!"
58 echo "    (or run ./toggle.sh in terminal)"
59 echo ""

```

## 10 docker-compose.yml

```

1 services:
2   warp:
3     image: qmcgaw/gluetun:v3.41.1
4     container_name: warp
5     hostname: nullexit-gateway
6     cap_add:
7       - NET_ADMIN
8     devices:
9       - /dev/net/tun:/dev/net/tun
10    ports:
11      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/udp" # AdGuard DNS: host binds 5354 (the mapped DNS-proxy
12        port); 5335 is AdGuard's in-container port
13      - "127.0.0.1:${DNS_PROXY_PORT:-5354}:5335/tcp" # AdGuard DNS (TCP)
14      - "41642:41642/udp" # Tailscale WireGuard (direct connection to host)
15      - "127.0.0.1:${SOCKS_PROXY_PORT:-1080}:1080/tcp" # SOCKS5 proxy (routed through WARP tunnel)
16      - "127.0.0.1:${TOR SOCKS PORT:-9050}:9050/tcp" # Tor SOCKS5 proxy (procedurally generated)
17      - "127.0.0.1:${TOR_CONTROL_PORT:-9051}:9051/tcp" # Tor Control port (procedurally generated)
18    environment:
19      - TZ=America/New_York
20      - VPN_SERVICE_PROVIDER=custom
21      - VPN_TYPE=wireguard
22      - WIREGUARD_PRIVATE_KEY=${WIREGUARD_PRIVATE_KEY}
23      - WIREGUARD_PUBLIC_KEY=${WIREGUARD_PUBLIC_KEY}
24      - WIREGUARD_ADDRESSES=${WIREGUARD_ADDRESSES}
25      - WIREGUARD_ENDPOINT_IP=${WARP_ENDPOINT_1:-162.159.192.1}
26      - WIREGUARD_ENDPOINT_PORT=2408
27      # Set to 25s (WireGuard default) to beat standard 30s hardware NAT/firewall UDP timeouts

```



```

27 - WIREGUARD_PERSISTENT_KEEPLIVE_INTERVAL=25s
28 # MTU Fix (Critical): Prevent packet fragmentation from double-encryption
29 - WIREGUARD_MTU=1280
30 # Give WARP more time to establish the connection before internal restart (default is 6s which is too
short)
31 - HEALTH_VPN_DURATION_INITIAL=30s
32 # Use plaintext DNS resolvers instead of DNS-over-TLS (DoT) to prevent startup failures on networks where
port 853 is blocked
33 - DNS_UPSTREAM_RESOLVER_TYPE=plain
34 - DNS_UPSTREAM_PLAIN_ADDRESSES=1.1.1.1:53,8.8.8.8:53
35 # Allow local network traffic to bypass WARP so Tailscale can establish fast, direct peer-to-peer
connections (no DERP relays) on the LAN
36 - FIREWALL_OUTBOUND_SUBNETS=192.168.0.0/16,172.16.0.0/12,10.0.0.0/8
37 # Built-in DNS blocking for ads, malicious domains, and tracking
38 # NOTE: This is the massive "Hagezi-style" built-in blocklist.
39 # You can turn these to 'on' for maximum security if you have spare RAM
40 # or are running this on a cloud exit node with adequate autonomous resources (requires approx 1-2GB of
RAM).
41 # Disabled intentionally: AdGuard Home provides the DNS filtering layer instead.
42 # If AdGuard fails, there is no fallback filtering, but the healthcheck
43 # dependency chain ensures Tailscale won't come up if AdGuard is down.
44 - BLOCK_MALICIOUS=off
45 - BLOCK_SURVEILLANCE=off
46 - BLOCK_ADS=off
47 dns:
48 - 1.1.1.1
49 sysctls:
50 - net.ipv4.ip_forward=1
51 volumes:
52 - ./post-rules.txt:/iptables/post-rules.txt:ro
53 restart: unless-stopped
54 healthcheck:
55   test: ["CMD-SHELL", "/gluetun-entrypoint healthcheck"]
56   interval: 2s
57   timeout: 5s
58   retries: 15
59   start_period: 60s
60 logging:
61   driver: "json-file"
62   options:
63     max-size: "10m"
64     max-file: "3"
65
66 tailscale:
67   image: tailscale/tailscale:v1.98.4
68   container_name: tailscale
69   network_mode: service:warp
70   depends_on:
71     warp:
72       condition: service_healthy
73     adguardhome:
74       condition: service_healthy
75   cap_add:
76     - NET_ADMIN
77     - NET_RAW
78     - SYS_MODULE
79   sysctls:
80     - net.ipv4.ip_forward=1
81     - net.ipv6.conf.all.forwarding=1
82   environment:
83     - TZ=America/New_York
84     - TS_EXTRA_ARGS=--advertise-exit-node
85     - PORT=41642
86     - TS_USERSPACE=false
87     - TS_AUTHKEY=${TS_AUTHKEY}
88     - TS_STATE_DIR=/var/lib/tailscale
89     # (Note: For headless/cloud deployments, you can bypass this footgun by
90     # using Tailscale Ephemeral Keys which auto-clean themselves).
91     - TS_AUTH_ONCE=true
92     # MTU Optimization: Force Tailscale packets to be smaller than WARP's 1280 MTU
93     # to prevent WireGuard-in-WireGuard fragmentation, completely eliminating bufferbloat
94     - TS_DEBUG_MTU=1200
95     # Enable Tailscale's built-in SOCKS5 proxy to route out of the WARP tunnel
96     - TS_SOCKS5_SERVER=:1080
97   devices:
98     - /dev/net/tun:/dev/net/tun
99   volumes:
100     - tailscale_state:/var/lib/tailscale
101   restart: unless-stopped
102   logging:
103     driver: "json-file"

```

```

104     options:
105         max-size: "10m"
106         max-file: "3"
107
108 routing-fix:
109     image: nullexit-routing-fix:v1.0.0
110     build:
111         context: ./scripts
112         dockerfile: Dockerfile.routing-fix
113     container_name: routing-fix
114     network_mode: service:warp
115     # NOTE: Deliberately NOT sharing tailscale's PID namespace the shared
116     # namespace caused routing-fix to be killed whenever tailscale restarted
117     # (e.g. during gluetun health-check flaps), which dropped all routing
118     # table changes and left traffic leaking through eth0 instead of tun0.
119     cap_add:
120         - NET_ADMIN
121         - SYSLOG
122     depends_on:
123         warp:
124             condition: service_healthy
125             # rule-compiler dependency removed: it is now a hash-gated one-off run by
126             # toggle.sh before `up` (15.8.8). This service boots on the existing
127             # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
128     environment:
129         - TZ=America/New_York
130         - WARP_ENDPOINT=${WARP_ENDPOINT_1:-162.159.192.1}
131         - BLOCKED_COUNTRIES=${BLOCKED_COUNTRIES}
132         - TAILSCALE_ALLOW_LAN_P2P=${TAILSCALE_ALLOW_LAN_P2P:-false}
133     volumes:
134         - ./scripts/routing-fix.sh:/routing-fix.sh:ro
135         # SECURITY: ./adguard/work/ is bind-mounted into this container.
136         # Never write to it from the host or another container. The
137         # single allowed writer is `rule-compiler` container manual
138         # edits corrupt the AdGuard cache + BoltDB session store.
139         - ./adguard/work/userfilters:/userfilters:ro
140         - ./adguard/work:/adguard_work:ro
141         - ./app
142     restart: unless-stopped
143     stop_grace_period: 1s
144     logging:
145         driver: "json-file"
146         options:
147             max-size: "1m"
148             max-file: "1"
149
150 adguardhome:
151     # WARNING: AdGuard is pre-configured with hardcoded credentials (admin/nullexit).
152     # This is fine for a local Mac running behind a NAT, but if you are deploying
153     # this on a public cloud VPS with port 80/3000 exposed, change this immediately!
154     image: adguard/adguardhome:v0.107.77
155     container_name: adguardhome
156     network_mode: service:warp
157     depends_on:
158         warp:
159             condition: service_healthy
160             # rule-compiler dependency removed: it is now a hash-gated one-off run by
161             # toggle.sh before `up` (15.8.8). This service boots on the existing
162             # compiled_rules.txt / ip_blocklist.ipset already present in ./adguard/work.
163     environment:
164         - TZ=America/New_York
165     healthcheck:
166         test: ["CMD-SHELL", "nc -z 127.0.0.1 5335 || exit 1"]
167         interval: 2s
168         timeout: 5s
169         retries: 10
170     volumes:
171         # SECURITY: ./adguard/work/ is bind-mounted into this container.
172         # Never write to it from the host or another container. The
173         # single allowed writer is `rule-compiler` container manual
174         # edits corrupt the AdGuard cache + BoltDB session store.
175         - ./adguard/work:/opt/adguardhome/work
176         - ./adguard/conf:/opt/adguardhome/conf
177     restart: unless-stopped
178     stop_grace_period: 2s
179     logging:
180         driver: "json-file"
181         options:
182             max-size: "10m"
183             max-file: "3"
184

```

```

185 rule-compiler:
186   image: nullexit-rule-compiler:v1.0.0
187   build:
188     context: ./scripts
189     dockerfile: Dockerfile.rule-compiler
190   container_name: rule-compiler
191   # `compile` profile: excluded from `docker compose up`. toggle.sh runs it
192   # on-demand via `docker compose run --build --rm rule-compiler`, hash-gated on
193   # black_list.txt/white_list.txt, so ~340k rules are re-deduped only when a list
194   # actually changes instead of every toggle (~28s saved on the common path, 15.8.8).
195   profiles: ["compile"]
196   volumes:
197     - ./app
198
199 tor:
200   build: ./docker/tor
201   image: nullexit-tor:v1.0.0
202   container_name: tor
203   network_mode: service:warp
204   depends_on:
205     warp:
206       condition: service_healthy
207   environment:
208     - TZ=America/New_York
209     - TOR_USE_BRIDGES=${TOR_USE_BRIDGES:-false}
210     - TOR_BRIDGE_LINES_FILE=${TOR_BRIDGE_LINES_FILE:-./tor-bridges.txt}
211     - TOR_PASSWORD=${ADGUARD_PASSWORD:-nullexit}
212     - TOR_EXCLUDE_EXIT_NODES=${TOR_EXCLUDE_EXIT_NODES:-{us}}
213   working_dir: /nullexit
214   volumes:
215     - ./:/nullexit:ro
216     - ./output.log:/output.log:rw
217   restart: unless-stopped
218   logging:
219     driver: "json-file"
220     options:
221       max-size: "10m"
222       max-file: "3"
223
224 volumes:
225   tailscale_state:
226
227 networks:
228   default:
229     ipam:
230       config:
231         - subnet: 10.200.1.0/24

```

## 11 scripts/common.sh

```

1 #!/bin/bash
2 # common.sh shared bash utilities for nullexit scripts
3
4 # Formatting
5 RED='\033[0;31m'
6 GREEN='\033[0;32m'
7 YELLOW='\033[1;33m'
8 BLUE='\033[0;34m'
9 BOLD='\033[1m'
10 NC='\033[0m'
11
12 # Constants & Environment
13 export MARKER_FILE="/tmp/nullexit-gateway-active.marker"
14 export PID_CAFFEINATE="/tmp/nullexit-cafeinate.pid"
15 export PID_DNS_WATCHER="/tmp/nullexit-dns-watcher.pid"
16 export DEFAULT_WARP_ENDPOINT_1="162.159.192.1"
17 export DEFAULT_WARP_ENDPOINT_2="162.159.193.1"
18
19 setup_standard_path() {
20   export PATH="/opt/homebrew/bin:/opt/homebrew/sbin:/usr/local/bin:$PATH"
21 }
22
23 step() { echo -e "\n[ $(date +%Y-%m-%d %H:%M:%S) ] ${BLUE}${BOLD} ${*}${NC}"; }
24 ok() { echo -e "[ $(date +%Y-%m-%d %H:%M:%S) ] ${GREEN} ${*}${NC}"; }
25 warn() { echo -e "[ $(date +%Y-%m-%d %H:%M:%S) ] ${YELLOW} ${*}${NC}"; }
26 fail() { echo -e "[ $(date +%Y-%m-%d %H:%M:%S) ] ${RED} ${*}${NC}"; }

```

```

27 die() { echo -e "\n[(date '+%Y-%m-%d %H:%M:%S')] ${RED} ${NC}\n"; exit 1; }
28
29 # Helper Functions
30
31 # Append a timestamped line to output.log (best-effort; never fails the caller).
32 # Used to record lifecycle breadcrumbs (EXEC/EXIT markers, phase changes) so a
33 # START/STOP that is killed or window-closed still leaves a retraceable trail.
34 log_line() {
35     echo "[(date '+%Y-%m-%d %H:%M:%S')] $" >> "${LOG_FILE:-output.log}" 2>/dev/null || true
36 }
37
38 # Run a lifecycle command with FULL observability: log the command, stream its
39 # combined stdout+stderr to the terminal AND append it to output.log via tee,
40 # then log the exit code explicitly. Returns the command's real exit code (not
41 # tee's). This closes the blind spot where `docker compose up` / `colima start`
42 # printed only to the terminal so when they fail (or the run is interrupted),
43 # output.log shows exactly what the last command was and whether it succeeded.
44 #
45 # Usage: run_logged docker compose up -d --build
46 run_logged() {
47     local logf="${LOG_FILE:-output.log}"
48     log_line "EXEC: $"
49     # PIPESTATUS[0] is the command's exit code; tee (PIPESTATUS[1]) is ignored.
50     "$@" 2>&1 | tee -a "$logf"
51     local rc=${PIPESTATUS[0]}
52     log_line "EXIT $rc: $"
53     return "$rc"
54 }
55
56 # Pure-bash timeout (no dependency on GNU coreutils' timeout)
57 # Runs a command with a safety cutoff so a wedged daemon can't hang the script.
58 run_with_timeout() {
59     local timeout_sec="$1"
60     shift
61     if [ $# -eq 0 ]; then return 1; fi
62
63     "$@" &
64     local cmd_pid=$!
65     CURRENT_BG_PID="$cmd_pid"
66
67     (
68         sleep "$timeout_sec"
69         if kill -0 "$cmd_pid" 2>> output.log; then
70             echo -e "\n[Timeout] Command '$*' exceeded $timeout_sec seconds. Terminating..." >&2
71             kill -15 "$cmd_pid" 2>> output.log || true
72             sleep 2
73             if kill -0 "$cmd_pid" 2>> output.log; then
74                 kill -9 "$cmd_pid" 2>> output.log || true
75             fi
76         fi
77     ) &
78     local watcher_pid=$!
79
80     # Disable set -e temporarily to capture exit status of wait
81     set +e
82     wait "$cmd_pid" 2>> output.log
83     local exit_status=$?
84     set -e
85
86     # Kill the watcher since the command finished
87     kill "$watcher_pid" 2>> output.log && wait "$watcher_pid" 2>> output.log || true
88
89     CURRENT_BG_PID=""
90     return "$exit_status"
91 }
92
93 # Start Colima and return the moment Docker is actually reachable, instead of
94 # blocking on colima 0.10.3's post-provision hang. Observed on macOS/vz: the VM
95 # reaches READY and creates the `colima` docker context in ~10-15s, then
96 # `colima start` frequently fails to return for ~110s (until a watchdog kills it)
97 # even though Docker is fully usable the entire time. We poll the colima docker
98 # context and proceed as soon as it answers, then reap the (usually hung) starter.
99 # The VM keeps running and `colima status` stays accurate. $@ = colima start args.
100 # Returns 0 once Docker responds, 1 if it never does within the cap.
101 colima_start_until_ready() {
102     local max_wait=90 waited=0 ready=false start_pid
103     echo "[(date '+%Y-%m-%d %H:%M:%S')] EXEC(bg): colima start $" >> output.log
104     colima start "$@" >> output.log 2>&1 &
105     start_pid=$!
106     while [ "$waited" -lt "$max_wait" ]; do
107         if run_with_timeout 5 docker --context colima info >/dev/null 2>&1; then

```

```

108     ready=true
109     break
110 fi
111 # If colima start exited on its own (genuine finish or hard failure), stop waiting.
112 if ! kill -0 "$start_pid" 2>/dev/null; then
113     run_with_timeout 5 docker --context colima info >/dev/null 2>&1 && ready=true
114     break
115 fi
116 sleep 3
117 waited=$((waited + 3))
118 done
119 # Reap the starter (harmless if it already exited); this does NOT stop the VM.
120 kill "$start_pid" 2>/dev/null || true
121 wait "$start_pid" 2>/dev/null || true
122 if [ "$ready" = "true" ]; then
123     docker context use colima >/dev/null 2>&1 || true
124     echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker ready after ~${waited}s (starter reaped)." >> output.log
125     return 0
126 fi
127 echo "[$(date '+%Y-%m-%d %H:%M:%S')] Colima Docker NOT ready after ${max_wait}s." >> output.log
128 return 1
129 }
130
131 # Hash the files that feed the long-running locally-built images that `docker
132 # compose up` starts (routing-fix, tor). toggle.sh uses this to skip
133 # `docker compose --build` when nothing changed: BuildKit re-evaluating those
134 # images on the 600MB Colima VM costs ~30s per toggle even when every layer is
135 # CACHED (15.8.8). Includes the Dockerfiles, everything they COPY (routing-fix.sh,
136 # tor/entrypoint.sh, the logger/ Go source), and docker-compose.yml. The
137 # rule-compiler image is NOT here it is a `compile`-profiled one-off, built+run
138 # by the gated compile step (15.8.8). Prints a sha256 hex digest, or empty.
139 compute_build_inputs_hash() {
140     local root="${SCRIPT_DIR:-.}" f
141     {
142         for f in \
143             "$root/scripts/Dockerfile.routing-fix" \
144             "$root/scripts/routing-fix.sh" \
145             "$root/docker/tor/Dockerfile" \
146             "$root/docker/tor/entrypoint.sh" \
147             "$root/docker-compose.yml"; do
148             printf ':%s: ' "$f"; cat "$f" 2>/dev/null
149         done
150         find "$root/scripts/logger" -type f 2>/dev/null \
151             | LC_ALL=C sort | while IFS= read -r f; do printf ':%s: ' "$f"; cat "$f" 2>/dev/null; done
152     } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
153 }
154
155 # Hash the ad-block lists the user controls: black_list.txt (custom blocks) and
156 # white_list.txt (the allow-list that governs what DNS resolves through). toggle.sh
157 # gates the ~28s rule recompile on this editing either list changes the digest
158 # and triggers a one-off recompile of compiled_rules.txt + ip_blocklist.ipset;
159 # an unchanged toggle skips it entirely (15.8.8). Prints a sha256 hex digest.
160 compute_rules_hash() {
161     local root="${SCRIPT_DIR:-.}"
162     { printf ':%s: ' "$root/black_list.txt" 2>/dev/null
163       printf ':%s: ' "$root/white_list.txt" 2>/dev/null
164     } | openssl dgst -sha256 2>/dev/null | awk '{print $NF}'
165 }
166
167 # Check if the gateway is active (either containers are running, or host DNS is hijacked)
168 is_gateway_active() {
169     # 1. Check if containers are running (suppress stderr to avoid errors if docker is down)
170     if run_with_timeout 15 docker compose ps --status running 2>/dev/null | grep -q 'warp'; then
171         echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (warp container is running)" >> "$LOG_FILE"
172         return 0
173     fi
174
175     # 2. Check if host DNS was hijacked (not 1.1.1.1 and not default/empty)
176     local current_dns=""
177     local active_svc=""
178     if [[ "$OSTYPE" == "darwin*" ]]; then
179         active_svc=$(get_active_service)
180         current_dns=$(networksetup -getdnsservers "$active_svc" 2>> output.log || true)
181     else
182         current_dns=$(resolvectl dns 2>/dev/null | awk '/Global|Link/ {for(i=4;i<=NF;i++) print $i}' | head -n1)
183     fi
184
185     if [[ -n "$current_dns" && "$current_dns" != "1.1.1.1" && ! "$current_dns" =~ "There aren't any DNS Servers"
186         ]]; then
187         echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (DNS was hijacked: $current_dns)" >> "$LOG_FILE"
188     fi
189 }

```

```

187     return 0
188 fi
189
190 # 3. Check if SOCKS proxy is enabled (macOS only)
191 if [[ "$OSTYPE" == "darwin"* ]]; then
192     local socks_proxy
193     socks_proxy=$(networksetup -getsocksfirewallproxy "$active_svc" 2>> output.log || true)
194     if echo "$socks_proxy" | grep -q "Enabled: Yes"; then
195         echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: TRUE (SOCKS proxy enabled)" >> "$LOG_FILE"
196         return 0
197     fi
198 fi
199
200 echo "[$(date '+%Y-%m-%d %H:%M:%S')] is_gateway_active: FALSE (Containers down, DNS clean, SOCKS disabled)"
201 >> "$LOG_FILE"
202 return 1
203 }
204
205 # Decide whether the gateway SHOULD be running, from persisted *intent* rather
206 # than a live health probe. This is the correct signal for toggle.sh's
207 # STOP-vs-START decision: a disable should tear down whatever the last START
208 # left behind it should NOT depend on whether Docker/Colima happens to be
209 # reachable right now.
210 #
211 # Why this exists: using is_gateway_active() for the decision means every toggle
212 # (including *disable*) runs `docker compose ps` first. When Colima is down the
213 # socket is absent, that call blocks for the full run_with_timeout window, and
214 # under `set -e` + the ERR trap a non-zero result at the `if then else fi`
215 # boundary could abort the script BEFORE the START body runs so the very path
216 # that would boot Colima never executed (observed as `EXIT: code=1 phase=start`).
217 #
218 # Signal source: MARKER_FILE is written at the end of a successful START
219 # (write_gateway_active_marker) and cleared at the top of STOP and in
220 # cleanup_handler. Reading it is instant and cannot hang or abort.
221 #
222 # ECHOES "running" or "stopped" to stdout (and ALWAYS returns 0). Callers read
223 # the string: [ "$(gateway_state)" = "running" ] &&
224 #
225 # CRITICAL why this echoes instead of using a 0/1 return code: on macOS
226 # /bin/bash 3.2, with `set -e` + an `ERR` trap + `set -x` (DEBUG_TRACE=true) all
227 # active, a function that does `return 1` triggers a SPURIOUS `set -e` abort in
228 # the caller even when the call is guarded (`if f; then`, or even `f || rc=$?`).
229 # The `return N` from the function, traced by `set -x`, trips the ERR trap. This
230 # is a genuine bash-3.2 interaction (see devref 15.11.8). Echoing a result and
231 # always returning 0 sidesteps `return`/`set -e`/`set -x` entirely the function
232 # can never contribute a non-zero status, so no caller can be aborted by it.
233 gateway_state() {
234     # Primary: persisted intent. Marker present => a START completed and no STOP
235     # has cleared it yet.
236     #
237     # NOTE: toggle.sh runs a "defensive stale-marker clear" near the top of every
238     # invocation (before this decision) to stop the wake-watcher firing against a
239     # crashed gateway. That would wipe the marker before we read it, so toggle.sh
240     # captures the marker's existence into GATEWAY_MARKER_AT_STARTUP *before* that
241     # clear and we honor it here. If the variable is unset (other callers, or the
242     # --restart pre-check which runs before the defensive clear), fall back to the
243     # live marker file.
244     if [ -n "${GATEWAY_MARKER_AT_STARTUP:-}" ]; then
245         if [ "$GATEWAY_MARKER_AT_STARTUP" = "yes" ]; then echo "running"; return 0; fi
246     elif [ -f "$MARKER_FILE" ]; then
247         echo "running"; return 0
248     fi
249
250     # Marker absent. Normally this means "stopped". But cover the edge case where
251     # the marker was lost (e.g. /tmp cleared, or gateway started by an older build)
252     # yet containers are genuinely up reconcile via is_gateway_active, but ONLY
253     # when Docker is actually reachable, so a dead-socket probe can never hang. On
254     # macOS Docker runs in the Colima VM; if that socket is absent, the gateway
255     # definitively is not running.
256     local docker_up="no"
257     if [[ "$OSTYPE" == "darwin"* ]]; then
258         if [ -S /var/run/docker.sock ] || [ -S "$HOME/.colima/default/docker.sock" ]; then docker_up="yes"; fi
259     else
260         if [ -S /var/run/docker.sock ]; then docker_up="yes"; fi
261     fi
262
263     if [ "$docker_up" = "yes" ] && is_gateway_active; then
264         echo "running"; return 0
265     fi
266     echo "stopped"
267     return 0

```



```

267 }
268
269 # Reads an environment variable from a file (default: .env), strips quotes and whitespace
270 read_env_var() {
271     local var_name="$1"
272     local env_file="{2:-.env}"
273     grep -E "^${var_name}=" "$env_file" 2>/dev/null | cut -d'=' -f2- | tr -d "\"\`" | sed -e 's/~[[:space:]]*~/'
274     -e 's/[[:space:]]*~/' || echo ""
275 }
276
277 # Load KILL_SWITCH variable from .env (defaults to false if not found/empty)
278 export KILL_SWITCH
279 KILL_SWITCH=$(read_env_var "KILL_SWITCH" 2>/dev/null | tr '[:upper:]' '[:lower:]')
280 if [ -z "$KILL_SWITCH" ]; then
281     KILL_SWITCH="false"
282 fi
283
284 # Function to get the active network service name (e.g., "Wi-Fi" or "USB 10/100 LAN")
285 get_active_service() {
286     if [[ "$OSTYPE" == "darwin"* ]]; then
287         local iface
288         iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
289
290         # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
291         if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
292             for i in en0 en1 en2 en3; do
293                 if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "
294                     inet "; then
295                     iface="$i"
296                     break
297                 fi
298             done
299         fi
300
301         # Default fallback if still empty or not enX
302         if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
303             iface="en0"
304         fi
305
306         # Map interface (e.g., en0) to service name (e.g., Wi-Fi)
307         local service
308         service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B 1 "Device: $iface" | head -n 1 |
309             sed -E 's/~([0-9\*]+\)\) //' )
310
311         if [ -n "$service" ]; then
312             echo "$service"
313         else
314             echo "Wi-Fi"
315         fi
316     else
317         # Get the interface used for default internet routing
318         local iface
319         iface=$(ip route get 1.1.1.1 2>/dev/null | awk '{print $5; exit}')
320         if [ -z "$iface" ]; then
321             echo ""
322             return 1
323         fi
324         echo "$iface"
325     fi
326 }
327
328 get_en0_service() {
329     local service
330     service=$(networksetup -listnetworkserviceorder 2>> output.log | grep -B1 "Device: en0" | head -1 | sed -E 's
331         /~([0-9\*]+\)\) //' || true)
332     if [ -z "$service" ]; then
333         echo "Wi-Fi"
334     else
335         echo "$service"
336     fi
337 }
338
339 get_warp_endpoint_1() { read_env_var "WARP_ENDPOINT_1" ".env" | grep -Eo '~[0-9\.]+' || echo "
340     $DEFAULT_WARP_ENDPOINT_1"; }
341 get_warp_endpoint_2() { read_env_var "WARP_ENDPOINT_2" ".env" | grep -Eo '~[0-9\.]+' || echo "
342     $DEFAULT_WARP_ENDPOINT_2"; }
343
344 restart_tailscaled_daemon() {
345     echo " [Tailscale Recovery] tailscaled daemon appears to be wedged/unresponsive."
346     echo " [Tailscale Recovery] Attempting to restart tailscaled service..."
347 }

```

```

342 if run_with_timeout 15 brew services restart tailscale >> output.log 2>&1; then
343     echo " [Tailscale Recovery] Successfully restarted tailscaled (user service)."
344 elif run_with_timeout 15 sudo -n brew services restart tailscale >> output.log 2>&1; then
345     echo " [Tailscale Recovery] Successfully restarted tailscaled (system service)."
346 else
347     echo " [Tailscale Recovery] WARNING: Failed to restart tailscaled. Manual intervention may be needed."
348 fi
349 sleep 3
350 }
351
352 disconnect_tailscale_host() {
353     local ts_bin="${1:-tailscale}"
354     if command -v "$ts_bin" >> output.log 2>&1; then
355         echo " Disconnecting host Tailscale from mesh..."
356
357         # Check if actually connected first (with timeout tailscaled could be wedged)
358         local status_ok=false
359         if run_with_timeout 5 "$ts_bin" status >> output.log 2>&1; then
360             status_ok=true
361         else
362             local status_exit=$?
363             # If it was a quick exit code 1 (disconnected), it is still reachable
364             if [ "$status_exit" -ne 143 ] && [ -S /var/run/tailscaled.socket ]; then
365                 status_ok=true
366             fi
367         fi
368
369         if [ "$status_ok" = "false" ]; then
370             restart_tailscaled_daemon
371         fi
372
373         # Explicitly reset any exit-node so it doesn't linger.
374         "$ts_bin" up >> output.log 2>&1 || true
375         if run_with_timeout 10 "$ts_bin" set --ssh=true --accept-dns=false --exit-node= >> output.log 2>&1; then
376             echo " [] Exit-node preference cleared."
377         else
378             echo " [!] tailscale up didn't respond (tailscaled may be wedged)"
379         fi
380
381         local ts_args=""
382         if is_kill_switch_enabled; then ts_args="--accept-risk=lose-ssh"; fi
383         if run_with_timeout 10 "$ts_bin" down $ts_args >> output.log 2>&1; then
384             echo " [] Tailscale disconnected."
385         else
386             echo " [!] tailscale down didn't respond"
387         fi
388     else
389         echo " [!] tailscale CLI not found skipping"
390     fi
391 }
392
393 stop_pidfile_daemon() {
394     local pid_file="$1"
395     local label="$2"
396     if [ -f "$pid_file" ]; then
397         local pid
398         pid=$(cat "$pid_file")
399         if kill -0 "$pid" 2>/dev/null; then
400             echo " Stopping $label (PID $pid)..."
401             sudo -n kill "$pid" 2>> output.log || kill "$pid" 2>> output.log || true
402             sleep 0.5
403             if kill -0 "$pid" 2>/dev/null; then
404                 sudo -n kill -9 "$pid" 2>> output.log || kill -9 "$pid" 2>> output.log || true
405             fi
406         fi
407         rm -f "$pid_file"
408     fi
409 }
410
411 disable_all_proxies() {
412     local svc="$1"
413     if [ -n "$svc" ]; then
414         sudo -n networksetup -setsocksfirewallproxystate "$svc" off 2>> output.log || true
415         sudo -n networksetup -setwebproxystate "$svc" off 2>> output.log || true
416         sudo -n networksetup -setsecurewebproxystate "$svc" off 2>> output.log || true
417     fi
418 }
419
420 flush_dns_cache() {
421     if [[ "$OSTYPE" == "darwin"* ]]; then
422         sudo -n dscacheutil -flushcache 2>> output.log || true

```

```

423     sudo -n killall -HUP mDNSResponder 2>> output.log || true
424 else
425     resolvectl flush-caches 2>> output.log || true
426 fi
427 }
428
429 wait_for_dhcp_settle() {
430     # Resolve physical interface (Wi-Fi or Ethernet)
431     local physical_iface
432     physical_iface=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: (Wi-Fi|Ethernet)/{
433         getline; print $2; exit}')
434     [ -z "$physical_iface" ] && physical_iface="en0"
435
436     echo -n "Waiting for physical network interface ($physical_iface) DHCP lease to settle..."
437     local settled=false
438     for attempt in {1..60}; do
439         PHYSICAL_GW=$(ipconfig getpacket "$physical_iface" 2>> output.log | awk -F'[]' '/router /{print $2}')
440         local local_ip
441         local_ip=$(ifconfig "$physical_iface" 2>/dev/null | awk '/inet /{print $2}')
442
443         if [ -n "$PHYSICAL_GW" ] && [[ "$PHYSICAL_GW" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
444             [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
445             [[ "$local_ip" != "169.254.*" ]]; then
446             echo "settled (Interface IP: $local_ip, Router IP: $PHYSICAL_GW)."
447             settled=true
448             break
449         fi
450
451         if [ "$attempt" -gt 15 ] && [ -n "$local_ip" ] && [[ "$local_ip" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]] && \
452             [[ "$local_ip" != "169.254.*" ]]; then
453             local fallback_gw
454             fallback_gw=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
455             if [ -n "$fallback_gw" ] && [[ "$fallback_gw" =~ ^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$ ]]; then
456                 PHYSICAL_GW="$fallback_gw"
457                 fi
458                 echo "static/settled detected (Interface IP: $local_ip, Gateway IP: ${PHYSICAL_GW:-unknown})."
459                 settled=true
460                 break
461             fi
462             echo -n "."
463             sleep 0.5
464         done
465         if [ "$settled" = "false" ]; then
466             echo "timed out or no active link. Proceeding anyway."
467         fi
468     }
469
470 reset_sharing_services() {
471     if [[ "$OSTYPE" == "darwin"* ]]; then
472         sudo -n killall sharingd rapportd 2>> output.log || true
473     fi
474 }
475
476 read_adguard_ip() {
477     if [ -f ".gateway_ip" ]; then
478         tr -d '\r' < ".gateway_ip" | awk 'NR==1{print $1;exit}'
479     fi
480 }
481
482 is_kill_switch_enabled() {
483     [ "$KILL_SWITCH" = "true" ]
484 }
485
486 add_warp_bypass_routes() {
487     local target="${1:-}"
488     local ep1
489     ep1=$(get_warp_endpoint_1)
490     local ep2
491     ep2=$(get_warp_endpoint_2)
492
493     if [[ "$OSTYPE" == "darwin"* ]]; then
494         local routing_arg=""
495         local msg_via=""
496         if [ -n "$target" ]; then
497             if [[ "$target" =~ ^en[0-9]+$ ]] || "$target" == "bridge*"; then
498                 routing_arg="-interface $target"
499                 msg_via="interface $target"
500             else
501                 routing_arg="$target"
502                 msg_via="gateway $target"
503             fi
504         fi
505     fi

```

```

502 else
503     local gateway_ip
504     gateway_ip=$(route get default 2>> output.log | awk '/gateway:/ {print $2}')
505     if [ -z "$gateway_ip" ]; then
506         local iface
507         iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
508         if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
509             iface="en0"
510         fi
511         routing_arg="-interface $iface"
512         msg_via="interface $iface"
513     else
514         routing_arg="$gateway_ip"
515         msg_via="gateway $gateway_ip"
516     fi
517 fi
518
519 echo -e "\nAdding host bypass routes for Cloudflare WARP endpoints via $msg_via..."
520 sudo -n route delete -host "$sep1" 2>/dev/null || true
521 sudo -n route delete -host "$sep2" 2>/dev/null || true
522 sudo -n route add -host "$sep1" $routing_arg >> output.log 2>&1 || true
523 sudo -n route add -host "$sep2" $routing_arg >> output.log 2>&1 || true
524
525 echo -e "Adding host bypass routes for Tailscale control plane via $msg_via..."
526 sudo -n route delete -net 192.200.0.0/24 2>/dev/null || true
527 sudo -n route add -net 192.200.0.0/24 $routing_arg >> output.log 2>&1 || true
528
529 echo -e "Adding host bypass routes for Tailscale DERP relays via $msg_via..."
530 local derp_ips
531 derp_ips=$(curl -s --connect-timeout 5 https://login.tailscale.com/derpmap/default | grep -oE '"IPv4"[[:space:]]*:[[:space:]]*[0-9.]+' | cut -d'"' -f4 | grep -E '^[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+$' || true)
532 if [ -n "$derp_ips" ]; then
533     echo "$derp_ips" > /tmp/nullexit-derp-ips.txt
534     for ip in $derp_ips; do
535         sudo -n route delete -host "$ip" 2>/dev/null || true
536         sudo -n route add -host "$ip" $routing_arg >> output.log 2>&1 || true
537     done
538 fi
539 else
540     local gateway_ip="${target}"
541     if [ -z "$gateway_ip" ]; then
542         gateway_ip=$(ip route show default 2>/dev/null | awk '/default/ {print $3}')
543     fi
544     if [ -n "$gateway_ip" ]; then
545         sudo ip route add "$sep1" via "$gateway_ip" >> output.log 2>&1 || true
546         sudo ip route add "$sep2" via "$gateway_ip" >> output.log 2>&1 || true
547     fi
548 fi
549 }
550
551 remove_warp_bypass_routes() {
552     local ep1
553     ep1=$(get_warp_endpoint_1)
554     local ep2
555     ep2=$(get_warp_endpoint_2)
556
557     if [[ "$OSTYPE" == "darwin*" ]]; then
558         echo -e "\nRemoving host bypass routes for Cloudflare WARP endpoints..."
559         sudo -n route delete -host "$sep1" >> output.log 2>&1 || true
560         sudo -n route delete -host "$sep2" >> output.log 2>&1 || true
561         echo -e "Removing host bypass routes for Tailscale control plane..."
562         sudo -n route delete -net 192.200.0.0/24 >> output.log 2>&1 || true
563         if [ -f /tmp/nullexit-derp-ips.txt ]; then
564             echo -e "Removing host bypass routes for Tailscale DERP relays..."
565             while read -r ip; do
566                 sudo -n route delete -host "$ip" >> output.log 2>&1 || true
567             done < /tmp/nullexit-derp-ips.txt
568             rm -f /tmp/nullexit-derp-ips.txt
569         fi
570     else
571         sudo ip route del "$sep1" >> output.log 2>&1 || true
572         sudo ip route del "$sep2" >> output.log 2>&1 || true
573     fi
574 }
575
576 bounce_wifi_interfaces() {
577     if [[ "$OSTYPE" == "darwin*" ]]; then
578         local wifi_port
579         wifi_port=$(networksetup -listallhardwareports 2>> output.log | awk '/Hardware Port: Wi-Fi/{getline; print $2}')
580         if [ -n "$wifi_port" ]; then

```

```

581     sudo -n networksetup -setairportpower "$wifi_port" off 2>> output.log || true
582     sleep 2
583     sudo -n networksetup -setairportpower "$wifi_port" on 2>> output.log || true
584     echo "Wi-Fi bounced (interface $wifi_port)"
585     sleep 3
586     else
587     echo "Could not detect Wi-Fi interface. Bouncing fallback interfaces..."
588     set +e
589     for iface in en0 en1 en2 en3 en4 en5; do
590         if ifconfig "$iface" >> output.log 2>&1; then
591             sudo -n ifconfig "$iface" down 2>> output.log || true
592             sudo -n ifconfig "$iface" up 2>> output.log || true
593         fi
594     done
595     set -e
596     echo "Fallback interfaces bounced"
597 fi
598 fi
599 }
600
601 get_active_interface() {
602     local iface
603     iface=$(route get default 2>> output.log | awk '/interface:/ {print $2}')
604
605     # If the default interface is empty or a VPN tunnel (like utunX), fallback to the active physical interface
606     if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
607         for i in en0 en1 en2 en3; do
608             if ifconfig "$i" 2>> output.log | grep -q "status: active" && ifconfig "$i" 2>> output.log | grep -q "
609                 inet "; then
610                 iface="$i"
611                 break
612             fi
613         done
614     fi
615
616     # Default fallback if still empty or not enX
617     if [[ -z "$iface" || ! "$iface" =~ ^en[0-9]+$ ]]; then
618         iface="en0"
619     fi
620     echo "$iface"
621 }
622
623 enable_killswitch() {
624     if [[ "$OSTYPE" == "darwin*" ]]; then
625         if ! is_kill_switch_enabled; then
626             return 0
627         fi
628
629         local ext_if
630         ext_if=$(get_active_interface)
631         local ep1
632         ep1=$(get_warp_endpoint_1)
633         local ep2
634         ep2=$(get_warp_endpoint_2)
635         local pf_conf_path
636         pf_conf_path="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)/pf.conf"
637
638         echo -e "\nEnabling macOS Packet Filter (PF) Kill-Switch on $ext_if..."
639
640         # Enable PF and capture the reference token
641         local token_out
642         token_out=$(sudo -n pfctl -E 2>> output.log || true)
643         local token
644         token=$(echo "$token_out" | awk '/Token/{print $NF}')
645         if [ -n "$token" ]; then
646             echo "$token" > /tmp/nullexit-pf-token.txt
647             echo " [] PF enabled with token $token."
648         else
649             echo " [!] WARNING: Could not capture PF reference token (may already be enabled or sudo failed)."
650         fi
651
652         # Load the rules into the anchor
653         if sudo -n pfctl -D "ext_if=$ext_if" -a com.apple/nullexit -f "$pf_conf_path" >> output.log 2>&1; then
654             echo " [] Ruleset loaded into anchor com.apple/nullexit."
655         else
656             echo " [!] ERROR: Failed to load ruleset into anchor."
657         fi
658
659         # Populate tables in the anchor
660         sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T flush >> output.log 2>&1 || true

```

```

661 sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$sep1" >> output.log 2>&1 || true
662 sudo -n pfctl -a com.apple/nullexit -t vpn_endpoints -T add "$sep2" >> output.log 2>&1 || true
663
664 if [ -f /tmp/nullexit-derp-ips.txt ]; then
665     local derp_ips
666     derp_ips=$(tr '\n' ' ' < /tmp/nullexit-derp-ips.txt)
667     if [ -n "$derp_ips" ]; then
668         sudo -n pfctl -a com.apple/nullexit -t derp_relays -T flush >> output.log 2>&1 || true
669         sudo -n pfctl -a com.apple/nullexit -t derp_relays -T add $derp_ips >> output.log 2>&1 || true
670     fi
671 fi
672 echo " [] Kill-switch tables populated."
673 fi
674 }
675
676 disable_killswitch() {
677     if [[ "$OSTYPE" == "darwin"* ]]; then
678         if ! is_kill_switch_enabled; then
679             return 0
680         fi
681         echo -e "\nDisabling macOS PF Kill-Switch..."
682
683         # Flush rules from anchor com.apple/nullexit
684         sudo -n pfctl -a com.apple/nullexit -F all >> output.log 2>&1 || true
685
686         # Release the enable reference if token is present
687         if [ -f /tmp/nullexit-pf-token.txt ]; then
688             local token
689             token=$(cat /tmp/nullexit-pf-token.txt)
690             if [ -n "$token" ]; then
691                 sudo -n pfctl -X "$token" >> output.log 2>&1 || true
692                 echo " [] Released PF enable reference token $token."
693             fi
694             rm -f /tmp/nullexit-pf-token.txt
695         else
696             echo " [] Rules flushed from anchor com.apple/nullexit."
697         fi
698     fi
699 }
700
701 write_host_ips() {
702     # Gather all active IPv4 addresses on the host and write to .host_ips
703     # routing-fix.sh will read this file to explicitly whitelist them for direct Tailscale P2P
704     local repo_dir
705     repo_dir="$(cd "$(dirname "${BASH_SOURCE[0]}")/.." && pwd)"
706
707     if [[ "$OSTYPE" == "darwin"* ]]; then
708         ifconfig | awk '/inet /{print $2}' | grep -v '127.0.0.1' > "$repo_dir/.host_ips" 2>/dev/null || true
709     else
710         ip -4 addr show 2>/dev/null | awk '/inet /{print $2}' | cut -d/ -f1 | grep -v '127.0.0.1' > "$repo_dir/.
711         host_ips" 2>/dev/null || true
712     fi
713 }
714
715 start_dns_watcher() {
716     if [[ "$OSTYPE" == "darwin"* ]]; then
717         local target_ip=$1
718         stop_pidfile_daemon "/tmp/nullexit-dns-watcher.pid" "background DNS Watcher"
719         echo " Starting background DNS Watcher for seamless Wi-Fi roaming..."
720         nohup bash -c "
721             source \"\$SCRIPT_DIR/scripts/common.sh\"
722             trap 'exit 0' SIGTERM SIGINT SIGHUP
723             while true; do
724                 ACTIVE_IF=$(get_active_service)
725                 if [ -n \"\$ACTIVE_IF\" ]; then
726                     CURRENT_DNS=$(networksetup -getdnsservers \"\$ACTIVE_IF\" 2>/dev/null)
727                     if [ \"\$CURRENT_DNS\" != \"\$target_ip\" ]; then
728                         networksetup -setdnsservers \"\$ACTIVE_IF\" \"\$target_ip\" >/dev/null 2>&1
729                     fi
730                 fi
731                 sleep 30
732             done
733             \" >> output.log 2>&1 &
734             echo $! > \"$DNS_WATCHER_PID_FILE"
735         else
736             local target_ip=$1
737             stop_dns_watcher
738             echo " Starting background DNS Watcher for seamless Wi-Fi roaming..."
739             nohup bash -c "
740                 trap 'exit 0' SIGTERM SIGINT SIGHUP

```



```

741     while true; do
742         if [ -n "\$ACTIVE_SERVICE\" ]; then
743             CURRENT_DNS=\$(resolvectl dns "\$ACTIVE_SERVICE\" 2>/dev/null | grep -oE
744             '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+' | head -1)
745             if [ "\$CURRENT_DNS\" != "\$target_ip\" ]; then
746                 resolvectl dns "\$ACTIVE_SERVICE\" "\$target_ip\" >/dev/null 2>&1 || true
747             fi
748             sleep 30
749         done
750         " >> output.log 2>&1 &
751         echo \$! > "\$DNS_WATCHER_PID_FILE"
752     fi
753 }
754
755 # SOCKS5 Proxy (WARP Tunnel)
756 # When Tailscale's exit node can't route through WARP (due to userspace
757 # forwarding), we use a SOCKS5 proxy running inside the warp container's
758 # network namespace. Connections created by this proxy go through the
759 # kernel routing table (table 200 -> tun0 -> WARP), unlike Tailscale's
760 # userspace exit node which bypasses it.
761 #
762 # The SOCKS5 proxy is exposed on localhost:\$SOCKS_PROXY_PORT via Docker port mapping.
763 # `networksetup -setsocksfirewallproxy` on macOS routes system TCP traffic
764 # through the proxy, which then goes through WARP.
765 #
766 # IMPORTANT: CLI tools like curl do NOT respect macOS system proxy settings.
767 # They must use --socks5-hostname or the ALL_PROXY env variable explicitly.
768
769 enable_socks_proxy() {
770     if [[ "\$OSTYPE" == "darwin*" ]]; then
771         local svc="\$1"
772         if [ -z "\$svc" ]; then
773             svc="\$ACTIVE_SERVICE"
774         fi
775
776         echo -n " Enabling system-wide SOCKS5 proxy on \$svc (localhost:\$SOCKS_PROXY_PORT)... "
777         sudo -n networksetup -setsocksfirewallproxy "\$svc" 127.0.0.1 \$SOCKS_PROXY_PORT 2>> output.log || true
778         sudo -n networksetup -setsocksfirewallproxystate "\$svc" on 2>> output.log || true
779
780         # Also set on en0 service for macOS resolver consistency
781         if [ "\$svc" != "\$ENO_SERVICE" ]; then
782             sudo -n networksetup -setsocksfirewallproxy "\$ENO_SERVICE" 127.0.0.1 \$SOCKS_PROXY_PORT 2>> output.log || true
783             sudo -n networksetup -setsocksfirewallproxystate "\$ENO_SERVICE" on 2>> output.log || true
784         fi
785
786         # IMPORTANT: Exclude local LAN and mDNS traffic from the proxy so P2P works natively
787         # without source-IP masking from the Colima VM bridge.
788         sudo -n networksetup -setproxybypassdomains "\$svc" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
789         172.16.0.0/12 "*.local" 2>> output.log || true
790         if [ "\$svc" != "\$ENO_SERVICE" ]; then
791             sudo -n networksetup -setproxybypassdomains "\$ENO_SERVICE" 127.0.0.1 localhost 192.168.0.0/16 10.0.0.0/8
792             172.16.0.0/12 "*.local" 2>> output.log || true
793         fi
794
795         echo "done."
796         echo " All TCP traffic now routed through gateway -> WARP tunnel -> internet."
797         echo " (CLI tools: export ALL_PROXY=socks5://127.0.0.1:\$SOCKS_PROXY_PORT)"
798     else
799         local svc="\$1"
800         echo " (Linux global SOCKS proxy configuration is desktop-environment specific, skipping.)"
801         echo " SOCKS5 proxy is active and listening on 127.0.0.1:\$SOCKS_PROXY_PORT"
802     fi
803 }
804
805 disable_socks_proxy() {
806     if [[ "\$OSTYPE" == "darwin*" ]]; then
807         for svc in "\$ACTIVE_SERVICE" "\$ENO_SERVICE"; do
808             sudo -n networksetup -setsocksfirewallproxystate "\$svc" off 2>> output.log || true
809         done
810         echo " SOCKS5 proxy disabled."
811     else
812         local svc="\$1"
813         echo " (Linux global SOCKS proxy configuration skipped.)"
814     fi
815 }
816
817 # Helper to restore host DNS to a clean state (used on script start, on failure,
818 # and after a successful STOP). Without this, a successful toggle-off leaves the
819 # host's resolver pinned to a now-dead gateway IP every lookup stalls for macOS's

```

```

818 # DNS timeout (~5s) before falling through to whatever next server is configured.
819 # With it, we always leave the host in `1.1.1.1 / no search domain`.
820 #
821 # We set DNS on BOTH the active service AND the en0 service because macOS's scutil DNS
822 # resolver is commonly scoped to en0 (Wi-Fi). A per-service change on e.g.
823 # "USB 10/100 LAN" is ignored by the system resolver unless the en0 service is also updated.
824 reset_dns() {
825     if [[ "$OSTYPE" == "darwin"* ]]; then
826         if [ -n "$ACTIVE_SERVICE" ]; then
827             for dns_svc in "$ACTIVE_SERVICE" "$ENO_SERVICE"; do
828                 networksetup -setsearchdomains "$dns_svc" "Empty" 2>> output.log || true
829                 networksetup -setdnsservers "$dns_svc" 1.1.1.1 2>> output.log || true
830             done
831         fi
832     else
833         if [ -n "$ACTIVE_SERVICE" ]; then
834             resolvectl domain "$ACTIVE_SERVICE" "" 2>/dev/null || true
835             resolvectl revert "$ACTIVE_SERVICE" 2>/dev/null || true
836         fi
837     fi
838 }
839
840 # Verify basic direct-internet connectivity after a recovery/teardown: DNS via
841 # 1.1.1.1, then HTTP (curl) or ping to 1.1.1.1. Sets the global INTERNET_OK to
842 # "true" on any successful check (callers read $INTERNET_OK in their summary).
843 # Always returns 0 so a bare call can't trip `set -e`. Used by recover.sh's macOS
844 # and Linux recovery paths (extracted from a formerly-duplicated block).
845 verify_internet_connectivity() {
846     step "Verifying internet connectivity"
847
848     # Wait a moment for the network to settle
849     sleep 2
850
851     INTERNET_OK=false
852
853     # Try DNS resolution first
854     if host -W 3 google.com 1.1.1.1 >> output.log 2>&1; then
855         ok "DNS resolution works (google.com via 1.1.1.1)"
856         INTERNET_OK=true
857     elif nslookup google.com 1.1.1.1 >> output.log 2>&1; then
858         ok "DNS resolution works (nslookup google.com via 1.1.1.1)"
859         INTERNET_OK=true
860     else
861         warn "DNS resolution check failed will still try ping/curl"
862     fi
863
864     # Try actual HTTP connectivity
865     if command -v curl >> output.log 2>&1; then
866         if curl -sf --max-time 5 https://1.1.1.1 >> output.log 2>&1; then
867             ok "Internet reachable via HTTP"
868             INTERNET_OK=true
869         elif curl -sf --max-time 5 http://1.1.1.1 >> output.log 2>&1; then
870             ok "Internet reachable via HTTP (plain)"
871             INTERNET_OK=true
872         else
873             warn "HTTP check failed"
874         fi
875     elif command -v ping >> output.log 2>&1; then
876         if ping -c 1 -W 3 1.1.1.1 >> output.log 2>&1; then
877             ok "Internet reachable via ping"
878             INTERNET_OK=true
879         else
880             warn "Ping check failed"
881         fi
882     fi
883
884     return 0
885 }
886
887 # Local DNS Proxy (Python)
888 # When the tailnet data plane cannot establish, we use a Python DNS proxy.
889 # It listens on UDP:53, receives DNS queries, forwards them over TCP to
890 # AdGuard at 127.0.0.1:${DNS_PROXY_PORT} (Docker port mapping), and returns the response.
891 # The Python script handles the DNS-over-TCP wire format (2-byte length prefix)
892 # correctly unlike socat which just forwards raw bytes.
893 #
894 # Python3 is built into macOS no external dependencies.
895 DNS_PROXY_PID=""
896
897 # Kill any leftover DNS proxy process
898 stop_local_dns_proxy() {

```

```

899 if [ -n "$DNS_PROXY_PID" ]; then
900     kill "$DNS_PROXY_PID" 2>/dev/null || true
901     wait "$DNS_PROXY_PID" 2>/dev/null || true
902     DNS_PROXY_PID=""
903 fi
904 # Clean up any stale Python DNS proxy processes
905 sudo -n pkill -f "dns-proxy.py" 2>/dev/null || true
906 }
907
908 # Start local DNS proxy (Python)
909 start_local_dns_proxy() {
910     if [ -z "$DNS_PROXY_BIN" ]; then
911         echo "Python not found. Python3 is built into macOS this shouldn't happen."
912         return 1
913     fi
914
915     if [ ! -f "$DNS_PROXY_SCRIPT" ]; then
916         echo "DNS proxy script not found at $DNS_PROXY_SCRIPT"
917         return 1
918     fi
919
920     # Kill any leftover process first
921     stop_local_dns_proxy
922
923     echo -n "Starting local DNS proxy via Python (UDP:53 TCP:${DNS_PROXY_PORT})... "
924     sudo -n bash -c "TARGET_PORT=$DNS_PROXY_PORT \"$DNS_PROXY_BIN\" \"$DNS_PROXY_SCRIPT\" \"\" &
925     DNS_PROXY_PID=$!
926     disown "$DNS_PROXY_PID" 2>/dev/null || true
927
928     # Give it a moment to bind
929     sleep 0.5
930
931     if kill -0 "$DNS_PROXY_PID" 2>/dev/null; then
932         echo "started (PID $DNS_PROXY_PID)."

```

```

980 # $1. Returns 0 iff the live state matches; non-zero if it can't be confirmed.
981 #
982 # This deliberately does NOT append a 1.1.1.1 fallback: macOS's resolver queries
983 # the list in order and falls back to the next entry on timeout, so on a
984 # misbehaving AdGuard (filter compile crash, OOM kill, etc.) the resolver still
985 # leaks to 1.1.1.1 and silently bypasses ad-blocking. With no fallback, a broken
986 # gateway manifests as a *visible* DNS outage that prompts the user to
987 # investigate which is the correct failure mode.
988 force_dns_to_gateway() {
989     if [[ "$OSTYPE" == "darwin"* ]]; then
990         local target_ip="$1"
991         if [ -z "$target_ip" ] || [ -z "$ACTIVE_SERVICE" ]; then
992             echo " [force_dns] ERROR: target_ip or ACTIVE_SERVICE is empty" >&2
993             return 1
994         fi
995
996         local attempt
997         for attempt in 1 2 3; do
998             echo -n " [force_dns] Attempt $attempt/3: setting DNS to $target_ip on \"$ACTIVE_SERVICE\" + \"
$ENO_SERVICE\"... "
999
1000             # Apply to BOTH the active service AND the en0 service the scutil DNS resolver
1001             # is scoped to en0 (Wi-Fi) and ignores per-service changes that don't include it.
1002             # Fail fast on ACTIVE_SERVICE; best-effort on ENO_SERVICE (it may not exist).
1003             networksetup -setsearchdomains "$ACTIVE_SERVICE" "ts.net" || true
1004             if ! networksetup -setdnsservers "$ACTIVE_SERVICE" "$target_ip"; then
1005                 echo "FAILED (networksetup error check permissions)"
1006                 sleep 1
1007                 continue
1008             fi
1009             networksetup -setsearchdomains "$ENO_SERVICE" "ts.net" || true
1010             networksetup -setdnsservers "$ENO_SERVICE" "$target_ip" || true
1011
1012             local entries
1013             entries=$(networksetup -getdnsservers "$ACTIVE_SERVICE" 2>> output.log \
1014                 | awk '/^(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\. (25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\
\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\.(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)$/ {print}')
1015             if echo "$entries" | grep -qFx "$target_ip"; then
1016                 echo "VERIFIED"
1017                 return 0
1018             fi
1019
1020             local current
1021             current=$(echo "$entries" | tr '\n' ' ' | sed 's/ $//')
1022             echo "NOT YET (current DNS: [${current:-none}])"
1023
1024             if [ "$attempt" = "3" ]; then sleep 2; else sleep 1; fi
1025         done
1026
1027         echo " [force_dns] FAILED after 3 attempts"
1028         return 1
1029     else
1030         local target_ip="$1"
1031
1032         for attempt in {1..3}; do
1033             echo -n " [force_dns] Attempt $attempt/3: setting DNS to $target_ip on $ACTIVE_SERVICE... "
1034             sudo -n resolvectl domain "$ACTIVE_SERVICE" "~ts.net" 2>/dev/null || true
1035             if ! sudo -n resolvectl dns "$ACTIVE_SERVICE" "$target_ip"; then
1036                 echo "FAILED (resolvectl error)"
1037             else
1038                 # Verify
1039                 entries=$(resolvectl dns "$ACTIVE_SERVICE" 2>> output.log | grep -oE '[0-9]+\.[0-9]+\.[0-9]+\.[0-9]+')
1040                 if [ "$entries" == "$target_ip" ]; then
1041                     echo "VERIFIED"
1042                     return 0
1043                 else
1044                     echo "MISMATCH (Resolver refused lock)"
1045                 fi
1046             fi
1047             sleep 2
1048         done
1049         return 1
1050     fi
1051 }

```

## 12 scripts/watcher.sh

```

1 #!/bin/bash
2 # scripts/watcher.sh nullexit post-wake / post-roam recovery watcher
3 #
4 # Long-running daemon. Launched by launchd as a LaunchAgent at
5 # ~/Library/LaunchAgents/com.nullexit.wake-recovery.plist
6 # (label com.nullexit.wake-recovery).
7 #
8 # Two independent listeners run as backgrounded pipelines:
9 #
10 # (1) WAKE listener `log stream` live-follows the macOS unified log for
11 # power-management events (lid closewake, manual wake, DarkWake from
12 # clamshell). Each event triggers run_recover (with a global debounce so
13 # a single wake that emits 2-3 log lines only causes ONE post-wake run).
14 #
15 # (2) NETWORK listener `scutil n.watch` on `State:/Network/Global/IPv4`
16 # fires on every Wi-Fi roam, ethernet swap, hotspot join, captive-portal
17 # re-bind, VPN up/down, etc. Same debounce.
18 #
19 # Both listeners call run_recover, which:
20 # - checks /tmp/nullexit-gateway-active.marker (only act when toggle.sh left
21 # the gateway up)
22 # - debounces (10s default) so multiple events don't pile up
23 # - shells out to `bash <repo>/recover.sh --post-wake` directly (no GUI auth prompt needed due to NOPASSWD
24 # sudoers)
25 #
26 # See devref.md 10.29 for the full Apple power management primer and the
27 # rationale behind using `log stream` + `scutil n.watch` from a shell daemon.
28 #
29 # NOTE: errexit (`set -e`) is intentionally NOT enabled. This is a long-running
30 # daemon, not a discrete-step script, and errexit actively breaks the reconnect
31 # loops below: any failed pipe (log stream on rotation, scutil on state churn)
32 # would kill the outer subshell before `echo "reconnecting..."; sleep 5;` runs,
33 # AND the parent's final `wait` would propagate a non-zero child exit and kill
34 # the entire watcher. Every error path is already wrapped in `|| true` /
35 # `2>/dev/null` fallbacks.
36 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
37 source "$SCRIPT_DIR/common.sh"
38 setup_standard_path
39 LOG="$SCRIPT_DIR/../output.log"
40 # Resolve our own directory; recover.sh lives one level up at the repo root.
41 # This makes the daemon install-agnostic no need to template the path
42 # in launchd, no need to keep a deploy-specific hardcoded location in sync.
43 RECOVER="$SCRIPT_DIR/../recover.sh"
44 MARKER="$MARKER_FILE"
45 DEBOUNCE_FILE="/tmp/nullexit-watcher.last-recovery"
46 DEBOUNCE_SECONDS="${NULLEXIT_DEBOUNCE_SECONDS:-10}"
47 # Post-recovery SETTLE cooldown. A post-wake run re-asserts the exit node
48 # (recover.sh: `tailscale set --exit-node`), which re-writes the host default
49 # route and that route IS State:/Network/Global/IPv4, the exact key Listener 2
50 # watches. The self-triggered change lands ~15-45s AFTER the run finishes, past
51 # the 10s debounce, so without a longer window the watcher re-fires post-wake in
52 # a ~40s self-feedback loop that only stops once routing happens to reach a fixed
53 # point (devref 15.12.21). This cooldown is armed ONLY after a recovery runs, so
54 # a genuine first wake/roam is still handled immediately; only the change we
55 # caused ourselves gets absorbed. A real new roam during the window is delayed at
56 # most this many seconds, which is acceptable.
57 COOLDOWN_FILE="/tmp/nullexit-watcher.settle-until"
58 POSTWAKE_SETTLE_SECONDS="${NULLEXIT_POSTWAKE_SETTLE_SECONDS:-45}"
59
60 exec >> "$LOG" 2>&1
61
62 # Single-instance lock
63 # Without this, repeated lid-closewake cycles under macOS Sonoma+ accumulate
64 # orphan watcher.sh processes: launchd suspends on sleep and SIGCONT-resumes
65 # on wake, and KeepAlive.Crashed=true caused relaunch on any non-zero exit;
66 # the listener grandchildren (log stream | while ... and (echo n.add ...;
67 # sleep 86400) | scutil | while ...) stay alive because pipe producers don't
68 # reliably respond to plain SIGTERM if the consumer is in a half-closed state.
69 #
70 # We solve it with a PID-file lock (POSIX-portable; macOS does NOT ship `flock`).
71 # The trap below removes the lock file on every exit so a stale lock never
72 # outlives a crashed watcher. On launch: read the existing PID, check it is
73 # alive AND is actually a `scripts/watcher.sh` process; if so, exit 0; else
74 # overwrite with our PID and proceed. The check-then-write window has a TOCTOU
75 # race but is microseconds wide and only matters for hand-launched duplicate
76 # test invocations; production is single-launch via launchctl load -w.
77 LOCK_FILE="/tmp/nullexit-watcher.lock"
78 LOCK_PID=""
79 if [ -e "$LOCK_FILE" ]; then

```

```

80 LOCK_PID=$(cat "$LOCK_FILE" 2>/dev/null || echo "")
81 if [ -n "$LOCK_PID" ] && kill -0 "$LOCK_PID" 2>/dev/null; then
82     # Verify the holder is actually a scripts/watcher.sh process (not some
83     # unrelated process that got the recycled PID). The exact-launchd pattern
84     # `bash <abs-path>/scripts/watcher.sh` is hard to accidentally match with
85     # a `grep scripts/watcher.sh .` invocation because we anchor on `^bash`
86     # followed by whitespace and end-of-line on the script path.
87     if ps -p "$LOCK_PID" -o args= 2>/dev/null | grep -qE "(^|[:space:])bash[:space:]+.*scripts/watcher\.sh$"
88     ; then
89         echo "[$(date -u +%FT%TZ)] watcher.sh: another instance holds $LOCK_FILE (pid=$LOCK_PID), exiting"
90         exit 0
91     fi
92 fi
93 echo $$ > "$LOCK_FILE"
94
95 echo "[$(date -u +%FT%TZ)] watcher.sh start (pid=$$ ppid=$PPID user=$(id -un) lock=acquired)"
96
97 # Treat launchd-resume as a wake signal. Apple's launchd SUSPENDS LaunchAgents
98 # during system sleep and restores them on wake; during that suspended gap,
99 # macOS's `log stream` cannot catch the wake event that prompted the resume.
100 # The next thing that happens after wake IS our own startup, so we always
101 # fire one post-wake run unconditionally here. The marker check + debounce
102 # in run_recover() make this idempotent: if the gateway isn't up or we just
103 # debounced, it no-ops. Without this, the FIRST wake-after-install goes
104 # unseen and the watcher appears broken until the SECOND wake.
105 # Helper: run recover.sh --post-wake if the gateway is active
106 run_recover() {
107     local why="$1"
108     if [ ! -f "$MARKER" ]; then
109         echo "[$(date -u +%FT%TZ)] $why no $MARKER, skip"
110         return 0
111     fi
112     local now last settle_until
113     now=$(date +%s)
114     # Post-recovery settle window: ignore the Global/IPv4 change our own last
115     # post-wake run caused (exit-node re-assert) so we don't self-feedback loop.
116     settle_until=$(cat "$COOLDOWN_FILE" 2>/dev/null || echo 0)
117     if [ "$now" -lt "$settle_until" ]; then
118         echo "[$(date -u +%FT%TZ)] $why settling ($(settle_until - now)s left in post-recovery cooldown), skip"
119         return 0
120     fi
121     last=$(cat "$DEBOUNCE_FILE" 2>/dev/null || echo 0)
122     if [ "$((now - last))" -lt "$DEBOUNCE_SECONDS" ]; then
123         echo "[$(date -u +%FT%TZ)] $why debounced ($(now - last)s since last, threshold ${DEBOUNCE_SECONDS}s)"
124         return 0
125     fi
126     echo "$now" > "$DEBOUNCE_FILE"
127     echo "[$(date -u +%FT%TZ)] $why bash $RECOVER --post-wake"
128     bash "$RECOVER" --post-wake
129     local exit_code=$?
130     now=$(date +%s)
131     echo "$now" > "$DEBOUNCE_FILE"
132     # Arm the settle cooldown so the route change this run just caused can't re-fire us.
133     echo "$((now + POSTWAKE_SETTLE_SECONDS))" > "$COOLDOWN_FILE"
134     echo "[$(date -u +%FT%TZ)] $why exit=$exit_code (settle until +${POSTWAKE_SETTLE_SECONDS}s, now=$now)"
135     return $exit_code
136 }
137
138 if [ -f "$MARKER" ]; then
139     run_recover "initial post-wake (launchd resume = wake signal)"
140 fi
141
142 # LAN P2P Auto-Detection
143 # Detects whether the current network allows safe Tailscale LAN P2P by:
144 # 1. Checking WPA2-Enterprise (802.1x) via system_profiler SPAirPortDataType always forces P2P=false
145 # because enterprise networks almost universally have AP Client Isolation.
146 # 2. AP Isolation probe sends a broadcast ping and checks if ANY LAN host responds.
147 # If no LAN hosts respond on a non-empty network, AP Isolation is active.
148 # Writes 'true' or 'false' to .lan_p2p_detected in the repo root.
149 # routing-fix.sh reads this file dynamically every 30s loop iteration.
150 P2P_OVERRIDE_FILE="$SCRIPT_DIR/../../lan_p2p_detected"
151
152 detect_lan_p2p_mode() {
153     local reason="unknown" allow="false"
154
155     # Step 1: Detect Wi-Fi security type via system_profiler (airport binary removed in macOS 14.4+)
156     # SPAirPortDataType lists all known networks; the first "Security:" line is the current association.
157     local security
158     security=$(system_profiler SPAirPortDataType 2>/dev/null \

```

```

160 | awk '/Current Network Information:/{found=1} found && /Security:/{print $NF; exit}'
161
162 if [ -z "$security" ]; then
163     # Not on Wi-Fi or system_profiler unavailable fail safe
164     reason="no-wifi" allow="false"
165 elif echo "$security" | grep -qi "Enterprise"; then
166     # WPA2/WPA3 Enterprise = 802.1x always AP isolated
167     reason="802.1x-enterprise" allow="false"
168 elif echo "$security" | grep -qi "Personal\|None\|Open"; then
169     # WPA2/WPA3 Personal or open probe for AP isolation
170     local gw
171     gw=$(route -n get 1.1.1.1 2>/dev/null | awk '/gateway:/{print $2}')
172     if [ -z "$gw" ]; then
173         reason="no-default-route" allow="false"
174     else
175         local responses
176         responses=$(ping -c 3 -t 1 -q "$gw" 2>/dev/null | awk '/packets transmitted/{print $4}')
177         if [ "${responses:-0}" -gt 0 ]; then
178             reason="wpa-personal-lan-reachable" allow="true"
179         else
180             reason="wpa-personal-ap-isolated" allow="false"
181         fi
182     fi
183 else
184     # Unknown security type fail safe
185     reason="unknown-security-type" allow="false"
186 fi
187
188 echo "[$(date -u +%FT%TZ)] P2P detect: security='${security:-none}' allow=${allow} (reason: ${reason})"
189 echo "$allow" > "$P2P_OVERRIDE_FILE"
190 write_host_ips
191 }
192
193 # Run once on startup
194 detect_lan_p2p_mode
195
196 # Listener 1: WAKE events via unified log
197 # Predicate picks up:
198 # * "didWake" regular kernel wake events (lid open, RTC alarm)
199 # * "Wake from Sleep" powerd-driven normal wake
200 # * "DarkWake" clamshell-display wake that didn't fully wake the
201 # Mac (relevant when lid opens during display sleep)
202 # * "Waking from" generic catch-all for variations in newer macOS
203 # `[c]` makes the match case-insensitive (the unified log uses mixed-case
204 # phrases). --info reduces noise by dropping debug/trace log levels.
205 # Subsystem-based predicate catches every powermanagement emission (willSleep,
206 # didWake, didDim, didUndim, DarkWake, etc.) regardless of message phrasing
207 # across kernel versions. Phrasing-based predicates ('didWake', 'Wake from Sleep')
208 # are fragile across macOS releases. We accept the extra log noise because the
209 # run_recover() debounce + marker check filter out anything that isn't a real
210 # gateway-relevant event.
211 (
212     # Reconnect loop: macOS rotates the unified log buffer periodically; when
213     # that happens, this `log stream` subscriber's pipe EOFs, the inner
214     # `while read` consumer terminates, and without this wrapper the listener
215     # would be silently dead for the rest of the watcher's lifetime.
216     while :; do
217         log stream --predicate \
218             'subsystem == "com.apple.powermanagement"' \
219             --style compact --info 2>/dev/null \
220             | while IFS= read -r line; do
221                 [ -z "$line" ] && continue
222                 run_recover "WAKE: $(echo "$line" | head -c 100)"
223             done
224         echo "[$(date -u +%FT%TZ)] log stream subshell exited; reconnecting in 5s"
225         sleep 5
226     done
227 ) &
228
229 # Listener 2: NETWORK state changes via scutil
230 # `scutil n.watch` on a state key prints SCEventUpdate lines whenever the key
231 # changes. We watch Global/IPv4 because that's the umbrella that catches link,
232 # address, route, and DNS changes for ALL interfaces in one namespace.
233 # The pipe is kept alive by an `sleep 86400` loop so scutil never sees EOF
234 # from us (which would silently terminate the watch).
235 (
236     # Reconnect loop: scutil can exit noisily on certain state changes (rare,
237     # but observed). Without the wrapper, the network listener's pipe EOFs and
238     # every subsequent roam goes unseen until launchd restarts us.
239     while :; do
240         (

```



```

241 # Watch both IPv4 and IPv6 umbrella state IPv6-only or dual-stack
242 # networks never fire on the IPv4 key alone, and the user's first roam
243 # in a hotel / airport / on cellular hotspot is often IPv6-first under
244 # NAT64.
245 echo "n.add State:/Network/Global/IPv4"
246 echo "n.add State:/Network/Global/IPv6"
247 echo "n.watch"
248 while true; do sleep 86400; done
249 ) | script -q /dev/null scutil 2>/dev/null \
250 | while IFS= read -r line; do
251     case "$line" in
252         *changedKey*|*State:/Network/Global/IPv*|*n.state*|*SCEventUpdate*)
253             detect_lan_p2p_mode
254             run_recover "NET: $(echo "$line" | head -c 100)"
255             ;;
256         *) ;;
257     esac
258 done
259 echo "[$(date -u +%FT%TZ)] scutil pipe exited; reconnecting in 5s"
260 sleep 5
261 done
262 ) &
263 # Listener 3: TUNNEL_FAILED_CLOSED.marker Watcher
264 # Polls for the fail-closed marker dropped by the routing-fix container.
265 (
266     last_notified=0
267     while ;; do
268         if [ -f "$SCRIPT_DIR/./TUNNEL_FAILED_CLOSED.marker" ]; then
269             now=$(date +%s)
270             # Notify once every 5 minutes (300s) if it stays failed
271             if [ "$(now - last_notified)" -ge 300 ]; then
272                 osascript -e 'display notification "VPN Tunnel Health Check Failed. Internet traffic is blackholed to
273 prevent IP leak." with title "NullExit Alert"' || true
274                 last_notified=$now
275             fi
276         else
277             last_notified=0
278         fi
279         sleep 3
280     done
281 ) &
282 # Lifetime
283 # `wait` blocks until both background pipelines exit (which they normally
284 # never do). launchctl `bootout`/SIGTERM triggers the trap, which kills the
285 # whole process group so the sleep-forever in the scutil pipe also dies.
286 # We catch TERM/INT/HUP for explicit signals and EXIT for any normal-exit path
287 # (so listeners always get cleaned up even if `exit 0` runs somewhere).
288 # `rm -f $LOCK_FILE` first so a stale lock can never block the next launch.
289 # `kill -TERM -P $$` then targets direct listener-children, then `kill 0`
290 # takes care of any backgrounded group-mates not reachable via the parent.
291 trap 'echo "[$(date -u +%FT%TZ)] watcher.sh terminating (signal/exit=$?)"; rm -f "$LOCK_FILE" 2>/dev/null ||
292 true; pkill -TERM -P $$ 2>/dev/null || true; kill 0 2>/dev/null || true; exit 0' TERM INT HUP EXIT
293 wait

```

## 13 scripts/crypto.sh

```

1 #!/usr/bin/env bash
2 # crypto.sh - Cryptographic Integrity Check
3
4 set -euo pipefail
5 cd "$(dirname "$0")/.."
6
7 if [ "$#" -ne 1 ] || { [ "$1" != "--sign" ] && [ "$1" != "--verify" ]; }; then
8     echo "Usage: $0 --sign | --verify"
9     exit 1
10 fi
11
12 if [ ! -f .env ]; then
13     echo "Error: .env not found."
14     exit 1
15 fi
16
17 SEED=$(grep "^NULLEXIT_SEED=" .env | cut -d '=' -f2)
18 if [ -z "$SEED" ]; then
19     echo "Error: NULLEXIT_SEED not found in .env."

```

```

20     exit 1
21 fi
22
23 if [ "$1" == "--sign" ]; then
24     echo "Generating cryptographic signatures using seed..."
25     rm -f .signatures
26     for file in toggle.sh recover.sh scripts/common.sh scripts/routing-fix.sh scripts/diagnose-host-leak.sh
27         scripts/sweep.sh scripts/watcher.sh scripts/pf.conf post-rules.txt docker-compose.yml black_list.txt
28         white_list.txt; do
29         if [ -f "$file" ]; then
30             hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
31             echo "file:$hash" >> .signatures
32             echo " [Signed] $file"
33         else
34             echo " [Skipped] $file not found"
35         fi
36     done
37     echo "Signatures saved to .signatures"
38     exit 0
39 fi
40
41 if [ "$1" == "--verify" ]; then
42     if [ ! -f .signatures ]; then
43         echo "Error: .signatures file missing. Run $0 --sign first."
44         exit 1
45     fi
46     FAIL=0
47     while IFS=: read -r file expected_hash; do
48         if [ -z "$file" ]; then continue; fi
49         if [ ! -f "$file" ]; then
50             echo "[FAIL] $file is missing!"
51             FAIL=1
52             continue
53         fi
54         actual_hash=$(openssl dgst -sha256 -hmac "$SEED" "$file" | awk '{print $NF}')
55         if [ "$actual_hash" != "$expected_hash" ]; then
56             echo "[FAIL] $file has been modified! Hash mismatch."
57             FAIL=1
58         fi
59     done < .signatures
60
61     if [ "$FAIL" -eq 1 ]; then
62         echo ""
63         echo "CRITICAL: Script integrity verification failed!"
64         echo "One or more core files have been modified without authorization."
65         echo "To authorize these changes, run: ./scripts/crypto.sh --sign"
66         echo ""
67         exit 1
68     fi
69     exit 0
70 fi

```

## 14 .aiignore

```

1 # Environment variables and secrets
2 .env
3 tor-bridges.txt

```

## 15 .claude/settings.local.json

```

1 {
2   "permissions": {
3     "allow": [
4       "Bash(grep *)",
5       "Bash(python3 /tmp/gen_nullexit_graph.py)",
6       "Bash(python3 -)",
7       "Bash(\\\"/Applications/Google Chrome.app/Contents/MacOS/Google Chrome\\\" --headless=new --disable-gpu --
8       window-size=1280,800 --virtual-time-budget=8000 --dump-dom \\\"file:///tmp/kg_debug.html\\\")",
9       "Bash(echo \\\"dangling-check-exit=?\\\")"
10    ]
11  }
12 }

```

## 16 .gitignore

```
1 # Environment variables and secrets
2 .env
3
4 # wgcf generated files containing keys and account info
5 wgcf-account.toml
6 wgcf-profile.conf
7
8 # Tailscale persistent state containing node identity
9 tailscale-state/
10
11 # Binaries
12 wgcf
13
14 # macOS App bundles
15 *.app/
16 *.app
17
18 # AdGuard Home persistent state
19 adguard/
20 .gateway_ip
21 logs.txt
22
23 #
24 stability_incident_report.md
25 output.log
26 .DS_Store
27 *.desktop
28 blocked.log
29 wtf.md
30 host-leak-diagnostic-*.txt
31 nullexit_review.md
32 nullexit_unified.tex
33 nullexit_unified.txt
34
35 nullexit_unified.synctex.gz
36 tor-bridges.txt
37
38 # Runtime state files
39 .host_ips
40 .lan_p2p_detected
41 .build_hash
42 .rules_hash
43 .signatures
44 TUNNEL_FAILED_CLOSED.marker
45
46 # Python cache
47 __pycache__/_
48 *.py[cod]
49
50 # LaTeX build artifacts
51 *.aux
52 *.log
53 *.out
54 *.toc
55 *.xdv
56 *.fdb_latexmk
57 *.fls
58 *.synctex*
59 refactor.md
60 sweep.md
61
62 # Obsidian Knowledge graph
63 nullexit-knowledge-graph/
```

## 17 Recover-Gateway.applescript

```
1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3     do shell script "true" with administrator privileges
4 on error
5     return
6 end try
7 tell application "Terminal"
```

```

8   activate
9   do script "cd " & quoted form of scriptDir & "; bash recover.sh"
10 end tell

```

## 18 Sweep-Gateway.applescript

```

1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3   do shell script "cd " & quoted form of scriptDir & " && ./scripts/crypto.sh --verify"
4 on error errMsg
5   display dialog "nullexit integrity check FAILED sweep will not run." & return & return & "A signed file (
      scripts/sweep.sh, scripts/common.sh, ) no longer matches .signatures. If you edited it on purpose, re-sign
      with:" & return & ".scripts/crypto.sh --sign" & return & return & errMsg buttons {"Abort"} default button
      "Abort" with icon stop
6   return
7 end try
8 tell application "Terminal"
9   activate
10  do script "cd " & quoted form of scriptDir & " && bash scripts/sweep.sh"
11 end tell

```

## 19 Toggle-Gateway.applescript

```

1 tell application "Finder" to set scriptDir to POSIX path of (container of (path to me) as alias)
2 try
3   do shell script "true" with administrator privileges
4 on error
5   return
6 end try
7 tell application "Terminal"
8   activate
9   do script "cd " & quoted form of scriptDir & "; ./toggle.sh"
10 end tell

```

## 20 benchmarks/benchmark.sh

```

1 #!/usr/bin/env bash
2
3 # benchmark.sh
4 # Automates the full benchmarking suite (both Python download test and Lighthouse HTML test)
5 # Runs the suite with Gateway ON, toggles OFF, runs again, then toggles back ON.
6
7 set -e
8
9 # Change to project root so we can access toggle.sh reliably
10 cd "$(dirname "$0")/.."
11
12 # Ensure jq is installed
13 if ! command -v jq &> /dev/null; then
14   echo "Error: 'jq' is not installed. Please run 'brew install jq'"
15   exit 1
16 fi
17
18 run_lighthouse() {
19   local url=$1
20   local file_prefix=$2
21   echo " [Lighthouse] Testing $url ..."
22
23   # Run lighthouse headlessly, silencing standard output
24   npx -y lighthouse "$url" --chrome-flags="--headless" --only-categories=performance --output=json --output-
     path="${file_prefix}.json" >/dev/null 2>&1
25
26   # Parse the results
27   local score=$(jq -r '.categories.performance.score' "${file_prefix}.json")
28   local tti=$(jq -r '.audits["interactive"].displayValue' "${file_prefix}.json")
29   local si=$(jq -r '.audits["speed-index"].displayValue' "${file_prefix}.json")
30
31   # Convert score to percentage

```

```

32     local score_pct=$(awk "BEGIN {print $score*100}")
33
34     echo "    -> Performance Score: ${score_pct}%"
35     echo "    -> Time to Interactive: $tti"
36     echo "    -> Speed Index: $si"
37
38     # Cleanup
39     rm -f "${file_prefix}.json"
40 }
41
42 echo "=====
43 echo "    PHASE 1: Testing with Gateway ON"
44 echo "=====
45 # 1. Run raw Python download test
46 python3 benchmarks/test_load.py
47
48 # 2. Run Lighthouse HTML render test
49 echo ""
50 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
51 run_lighthouse "https://www.cnet.com/" "on_cnet"
52 run_lighthouse "https://www.independent.co.uk/" "on_ind"
53
54
55 echo ""
56 echo "=====
57 echo "    PHASE 2: Toggling Gateway OFF"
58 echo "=====
59 ./toggle.sh
60 echo "Waiting 10 seconds for macOS Wi-Fi to stabilize after DNS reset..."
61 sleep 10
62
63
64 echo ""
65 echo "=====
66 echo "    PHASE 3: Testing with Gateway OFF"
67 echo "=====
68 # 1. Run raw Python download test
69 python3 benchmarks/test_load.py
70
71 # 2. Run Lighthouse HTML render test
72 echo ""
73 echo "Starting Lighthouse HTML Tests (this takes ~30-60s per site)..."
74 run_lighthouse "https://www.cnet.com/" "off_cnet"
75 run_lighthouse "https://www.independent.co.uk/" "off_ind"
76
77
78 echo ""
79 echo "=====
80 echo "    PHASE 4: Restoring Gateway ON"
81 echo "=====
82 ./toggle.sh
83
84 echo ""
85 echo "=====
86 echo " BENCHMARKS COMPLETE!"
87 echo "You can scroll up to compare Phase 1 (AdGuard ON) vs Phase 3 (AdGuard OFF)."
88 echo "=====

```

## 21 benchmarks/test\_load.py

```

1 import urllib.request
2 import urllib.error
3 import time
4 from html.parser import HTMLParser
5 import concurrent.futures
6
7 class ResourceParser(HTMLParser):
8     def __init__(self):
9         super().__init__()
10        self.urls = set()
11
12    def handle_starttag(self, tag, attrs):
13        attrs = dict(attrs)
14        url = None
15        if tag in ['img', 'script']:
16            url = attrs.get('src')

```

```

17         elif tag == 'link' and attrs.get('rel') == 'stylesheet':
18             url = attrs.get('href')
19
20         if url and url.startswith('http'):
21             self.urls.add(url)
22
23     def fetch_url(url):
24         try:
25             req = urllib.request.Request(url, headers={'User-Agent': 'Mozilla/5.0'})
26             urllib.request.urlopen(req, timeout=3)
27             return True
28         except Exception:
29             return False
30
31     def measure_url(target_url):
32         print(f"\nTesting {target_url} Load...")
33         start_total = time.time()
34
35         # Fetch main HTML
36         try:
37             req = urllib.request.Request(target_url, headers={'User-Agent': 'Mozilla/5.0 (Macintosh; Intel Mac OS X
38             10_15_7) AppleWebKit/537.36'})
39             response = urllib.request.urlopen(req, timeout=5)
40             html = response.read().decode('utf-8', errors='ignore')
41         except Exception as e:
42             print(f"Failed to fetch {target_url}: {e}")
43             return
44
45         parser = ResourceParser()
46         parser.feed(html)
47         urls = list(parser.urls)
48         print(f"Found {len(urls)} subresources to fetch (scripts, images, css)...")
49
50         # Concurrently fetch subresources
51         success_count = 0
52         fail_count = 0
53         max_threads = min(64, max(1, len(urls)))
54         with concurrent.futures.ThreadPoolExecutor(max_workers=max_threads) as executor:
55             results = list(executor.map(fetch_url, urls))
56
57         success_count = sum(1 for r in results if r)
58         fail_count = len(results) - success_count
59
60         total_time = time.time() - start_total
61         print(f"Time taken: {total_time:.2f} seconds")
62         print(f"Successfully fetched: {success_count}, Failed/Blocked: {fail_count}")
63
64     urls_to_test = [
65         "https://www.dailymail.co.uk/home/index.html",
66         "https://www.cnet.com/",
67         "https://www.independent.co.uk/"
68     ]
69
70     for u in urls_to_test:
71         measure_url(u)

```

## 22 black\_list.txt

```

1 ad.doubleclick.net
2 adclick.g.doubleclick.net
3 adeventtracker.spotify.com
4 adfstat.yandex.ru
5 ads-api.tiktok.com
6 ads-api.twitter.com
7 ads-fa.spotify.com
8 ads-sg.tiktok.com
9 ads.tiktok.com
10 adserver.unityads.unity3d.com
11 adsfs.oppomobile.com
12 an.facebook.com$important
13 connect.facebook.net$important
14 pixel.facebook.com$important
15 analytics.query.yahoo.com
16 analytics.s3.amazonaws.com
17 analytics.spotify.com
18 analyticsengine.s3.amazonaws.com

```

```

19 appmetrica.yandex.ru
20 audio2.spotify.com
21 bdapi-ads.realmemobile.com
22 bdapi-in-ads.realmemobile.com
23 bounceexchange.com
24 bs.serving-sys.com
25 business-api.tiktok.com
26 ck.ads.oppomobile.com
27 crashdump.spotify.com
28 data.ads.oppomobile.com
29 data.mistat.india.xiaomi.com
30 data.mistat.rus.xiaomi.com
31 data.mistat.xiaomi.com
32 desktop.spotify.com
33 doubleclick.net
34 ds.metric.spotify.com
35 gemini.yahoo.com
36 googleads.g.doubleclick.net
37 googleadservices.com
38 googlesyndication.com
39 grs.hicloud.com
40 gtm.spotify.com
41 iot-eu-logser.realme.com
42 iot-logser.realme.com
43 log.byteoversea.com
44 log.fc.yahoo.com
45 log.spotify.com
46 m.doubleclick.net
47 mediavisor.doubleclick.net
48 metrics.spotify.com
49 metrika.yandex.ru
50 omaze.com
51 pagead2.googlesyndication.com
52 securepubads.g.doubleclick.net
53 static.doubleclick.net
54 stats.g.doubleclick.net
55 tracking.rus.miui.com
56 udcml.yahoo.com
57 v.jwpcdn.com
58 weblb-wg.gslb.spotify.com
59 webview.unityads.unity3d.com
60 diagassets.apple.com
61 iphonesubmissions.apple.com
62 radarsubmissions.apple.com
63 metrics.apple.com
64 xp.apple.com
65 heads-fa.spotify.com
66 heads4.spotify.com
67 pixel.spotify.com
68 segs.spotify.com
69 audio-fa.spotify.com
70 events.data.microsoft.com
71 datarouter-prod.ak.epicgames.com
72 nel.cloudflare.com
73 eu-mobile.events.data.microsoft.com
74 stocks-analytics-events.apple.com

```

## 23 cloud-init.yaml

```

1 #cloud-config
2 # Generic Cloud-Init Bootstrap Script for Remote Exit Nodes (Ubuntu/Debian)
3 # Paste this into the "User Data" or "Initialization Script" section when creating a VPS.
4
5 package_update: true
6 package_upgrade: true
7
8 packages:
9   - apt-transport-https
10   - ca-certificates
11   - curl
12   - gnupg
13   - lsb-release
14   - git
15   - python3
16   - python3-pip
17

```



```

18 runcmd:
19 # 1. Install Docker natively
20 - curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-
  archive-keyring.gpg
21 - echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg]
  https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/
  docker.list > /dev/null
22 - sudo apt-get update
23 - sudo apt-get install -y docker-ce docker-ce-cli containerd.io docker-compose-plugin
24
25 # 2. Add default user (ubuntu) to docker group
26 - usermod -aG docker ubuntu || true
27 - usermod -aG docker debian || true
28
29 # 3. Create the installation directory
30 - mkdir -p /opt/nullexit
31 - chown -R ubuntu:ubuntu /opt/nullexit || chown -R debian:debian /opt/nullexit
32
33 # 4. Enable IPv4 and IPv6 forwarding for VPN routing
34 - echo "net.ipv4.ip_forward=1" >> /etc/sysctl.d/99-tailscale.conf
35 - echo "net.ipv6.conf.all.forwarding=1" >> /etc/sysctl.d/99-tailscale.conf
36 - sysctl -p /etc/sysctl.d/99-tailscale.conf
37
38 final_message: "The remote exit node has been fully bootstrapped. SSH in, clone the repo to /opt/nullexit, add
  your .env, and run ./setup-linux.sh!"

```

## 24 docker/tor/Dockerfile

```

1 FROM debian:bookworm-slim
2
3 RUN apt-get update && \
4     apt-get install -y tor obfs4proxy gosu bash && \
5     rm -rf /var/lib/apt/lists/*
6
7 COPY entrypoint.sh /entrypoint.sh
8 RUN chmod +x /entrypoint.sh
9
10 ENTRYPOINT ["/entrypoint.sh"]

```

## 25 docker/tor/entrypoint.sh

```

1 #!/bin/bash
2 set -e
3
4 TORRC="/etc/tor/torrc"
5 mkdir -p /etc/tor /var/lib/tor
6
7 # Base configuration: SOCKS 9050, Control 9051, TransPort 9040, DNSPort 5353 all bound to
8 # 0.0.0.0 (required for Docker/Colima host port-mapping). The ControlPort is secured via a
9 # dynamic HashedControlPassword (added below when TOR_PASSWORD is set); see devref.md 15.10.2.
10 cat <<EOF > $TORRC
11 SocksPort 0.0.0.0:9050
12 ControlPort 0.0.0.0:9051
13 CookieAuthentication 0
14 Log notice stdout
15 DataDirectory /var/lib/tor
16
17 # Transparent proxying & DNS resolution for .onion mapping
18 TransPort 0.0.0.0:9040
19 DNSPort 0.0.0.0:5353
20 AutomapHostsOnResolve 1
21 VirtualAddrNetworkIPv4 198.18.0.0/15
22 EOF
23
24 if [ -n "$TOR_PASSWORD" ]; then
25     HASH=$(tor --hash-password "$TOR_PASSWORD" | tail -n 1)
26     echo "HashedControlPassword $HASH" >> $TORRC
27 fi
28
29 if [ -n "$TOR_EXCLUDE_EXIT_NODES" ]; then
30     echo "ExcludeExitNodes $TOR_EXCLUDE_EXIT_NODES" >> $TORRC
31     echo "StrictNodes 1" >> $TORRC

```

```

32 fi
33
34 if [ "${TOR_USE_BRIDGES:-false}" = "true" ]; then
35     BRIDGE_FILE="${TOR_BRIDGE_LINES_FILE:-/tor-bridges.txt}"
36
37     if ! grep -vE "^[[:space:]]*#" "$BRIDGE_FILE" | grep -q '^[[:space:]]'; then
38         MSG="[$(date '+%Y-%m-%d %H:%M:%S')] [CRITICAL] [Tor Proxy] TOR_USE_BRIDGES is enabled, but no bridge
39         lines found at $BRIDGE_FILE. Proxy startup ABORTED. Get bridges from https://bridges.torproject.org."
40         echo "$MSG"
41         if [ -w "/output.log" ]; then
42             echo "$MSG" >> /output.log
43         fi
44         # Sleep indefinitely to prevent docker-compose 'unless-stopped' from triggering a crash loop
45         exec sleep infinity
46     fi
47
48     echo "UseBridges 1" >> $TORRC
49     echo "ClientTransportPlugin obfs4 exec /usr/bin/obfs4proxy" >> $TORRC
50
51     while IFS= read -r line; do
52         # Ignore comments and empty lines
53         [[ -z "$line" || "$line" == \#* ]] && continue
54         echo "Bridge $line" >> $TORRC
55     done < "$BRIDGE_FILE"
56 fi
57
58 # Debian tor package uses 'debian-tor' user
59 chown -R debian-tor:debian-tor /var/lib/tor /etc/tor
60
61 exec gosu debian-tor /usr/bin/tor -f $TORRC

```

## 26 launchd/com.nullexit.wake-recovery.plist

```

1 <?xml version="1.0" encoding="UTF-8"?>
2 <!DOCTYPE plist PUBLIC "-//Apple//DTD PLIST 1.0//EN" "http://www.apple.com/DTDs/PropertyList-1.0.dtd">
3 <plist version="1.0">
4 <dict>
5     <key>Label</key>
6     <string>com.nullexit.wake-recovery</string>
7
8     <!-- WorkingDirectory + relative program arg keeps this plist install-agnostic.
9     The source-of-truth for the install path is __NULLEXIT_HOME__, which
10    setup.sh or the user must substitute with their absolute nullexit
11    install root at install time. Program arg is `scripts/watcher.sh`
12    RELATIVE to WorkingDirectory, so launchd's evaluated cwd IS the
13    install root and BASH_SOURCE inside scripts/watcher.sh still
14    resolves to the real <install>/scripts/watcher.sh, preserving the
15    SCRIPT_DIR pattern. After install-time `sed -i ' ' 's|__NULLEXIT_HOME__|<abs_path>|'`,
16    the plist is portable to any clone location. -->
17     <key>WorkingDirectory</key>
18     <string>__NULLEXIT_HOME__</string>
19
20     <key>ProgramArguments</key>
21     <array>
22         <string>/bin/bash</string>
23         <string>scripts/watcher.sh</string>
24     </array>
25
26     <!-- Start the daemon immediately on launchctl load (instead of waiting
27     for the next user login). Combined with keepAlive below this means
28     once installed: load once, runs forever across logouts/reboots. -->
29     <key>RunAtLoad</key>
30     <true/>
31
32     <!-- Restart policy:
33     SuccessfulExit=false    SIGTERM / clean exit does NOT relaunch (so
34                             `launchctl bootout` cleanly stops, not loops).
35     Crashed=false          HARD crashes ALSO do NOT relaunch. We've seen
36                             that repeated sleep/wake cycles on macOS
37                             Sonoma+ accumulate orphan watcher.sh PIDs
38                             because SIGCONT/resume doesn't reliably kill
39                             the prior instance. The watcher.sh script
40                             now enforces single-instance via a manual
41                             PID-file and process-identity check at the
42                             top, so a re-launch on crash is
43                             actively dangerous (it would skip the lock

```

```

44                                     and exit; better to surface the crash in
45                                     logs than to silently 0-process). -->
46
47     <key>KeepAlive</key>
48     <dict>
49         <key>SuccessfulExit</key>
50         <false/>
51         <key>Crashed</key>
52         <false/>
53     </dict>
54
55     <!-- Background hint to App Nap so it doesn't get suspended when
56         the user is inactive. We're the entire reason this LaunchAgent
57         needs to NOT nap. -->
58     <key>ProcessType</key>
59     <string>Background</string>
60
61     <!-- launchd-level stdout/stderr (the bash launcher shell). The script
62         itself writes to /tmp/nullexit-watcher.log via exec >>. These
63         separate channels help us tell "launchd crashed" from "the bash
64         wrapper piped to recover.sh errored". -->
65     <key>StandardOutPath</key>
66     <string>/tmp/nullexit-watcher.out.log</string>
67     <key>StandardErrorPath</key>
68     <string>/tmp/nullexit-watcher.err.log</string>
69 </dict>
</plist>

```

## 27 post-rules.txt

```

1 iptables -A FORWARD -i tailscale0 -o tun0 -j ACCEPT
2 iptables -A FORWARD -i tun0 -o tailscale0 -j ACCEPT
3 iptables -A INPUT -i tailscale0 -j ACCEPT
4 iptables -A INPUT -i eth0 -p udp --dport 41642 -j ACCEPT
5 iptables -t nat -A POSTROUTING -o tun0 -j MASQUERADE
6 # Clamp MSS to 1120. With Tailscale overhead on top of WARP (each WireGuard encapsulation adds ~60 bytes),
7 # the safe MSS for double-tunneled traffic is 1120 to prevent strict-MSS endpoints from stalling.
8 # NOTE: this 1120 integer is rewritten at boot by toggle.sh from GATEWAY_MSS in .env
9 # (Gluetun's parser can't interpolate variables) change GATEWAY_MSS in .env, not this line.
10 iptables -t mangle -A POSTROUTING -p tcp --tcp-flags SYN,RST SYN -j TCPMSS --set-mss 1120
11
12 iptables -t nat -A PREROUTING -i tailscale0 -p udp --dport 53 -j REDIRECT --to-port 5335
13 iptables -t nat -A PREROUTING -i tailscale0 -p tcp --dport 53 -j REDIRECT --to-port 5335
14
15 # Also redirect DNS from Docker bridge interface so the host can reach AdGuard via localhost:53
16 iptables -t nat -A PREROUTING -i eth0 -p udp --dport 53 -j REDIRECT --to-port 5335
17 iptables -t nat -A PREROUTING -i eth0 -p tcp --dport 53 -j REDIRECT --to-port 5335
18
19 # Block IPv6 exit-node traffic to prevent leakage (Gluetun/WARP is IPv4-only)
20 ip6tables -A FORWARD -i tailscale0 -j DROP
21
22 # Block DNS-over-TLS (Port 853) to force Android Private DNS to fallback to Port 53
23 iptables -A FORWARD -i tailscale0 -p tcp --dport 853 -j REJECT --reject-with tcp-reset
24 iptables -A FORWARD -i tailscale0 -p udp --dport 853 -j REJECT
25
26 # Block QUIC (UDP 443) to force Facebook/Google to fallback to TCP.
27 # UDP bypasses TCP MSS clamping, causing fragmentation. Over weak Wi-Fi (far from gateway),
28 # dropping a single fragment destroys the entire packet, stalling image downloads.
29 iptables -A FORWARD -i tailscale0 -p udp --dport 443 -j REJECT

```

## 28 scripts/Dockerfile.routing-fix

```

1 FROM golang:1.22-alpine AS builder
2 WORKDIR /app
3 COPY logger/ ./logger/
4 RUN cd logger && go build -o logger main.go
5
6 FROM alpine:3.20
7 RUN apk add --no-cache iptables iproute2 curl ipset
8 COPY --from=builder /app/logger/logger /usr/local/bin/nullexit-logger
9 COPY routing-fix.sh /routing-fix.sh
10 CMD ["sh", "/routing-fix.sh"]

```

## 29 scripts/Dockerfile.rule-compiler

```
1 FROM golang:1.22-alpine AS builder
2 WORKDIR /app
3 COPY rule-compiler/ ./rule-compiler/
4 RUN cd rule-compiler && go build -o sync-rules main.go
5
6 FROM alpine:3.20
7 COPY --from=builder /app/rule-compiler/sync-rules /usr/local/bin/sync-rules
8 WORKDIR /app
9 CMD ["sync-rules"]
```

## 30 scripts/Toggle-Gateway.bat

```
1 <# :
2 @echo off
3 powershell -NoProfile -ExecutionPolicy Bypass -Command "Invoke-Expression ([System.IO.File]::ReadAllText('%-f0
4 ') -replace '^(?s).*?#\s*>\s*\r?\n','')
5 exit /b
6 #>
7 # Toggle-Gateway.bat (Hybrid Batch-PowerShell Toggler) nullexit Gateway Toggler for Windows
8 # Mirrors the sophisticated error handling, checks, and automation of toggle.sh
9
10 # 1. Administrator Check & Elevation
11 $isAdmin = ([Security.Principal.WindowsPrincipal][Security.Principal.WindowsIdentity]::GetCurrent()).IsInRole([
12 Security.Principal.WindowsBuiltInRole]::Administrator)
13 if (-not $isAdmin) {
14 Write-Host "Requesting Administrator privileges..." -ForegroundColor Yellow
15 Start-Process PowerShell -ArgumentList "-NoProfile -ExecutionPolicy Bypass -Command `\"Invoke-Expression ([
16 System.IO.File]::ReadAllText('$PSCmdPath') -replace '^(?s).*?#\s*>\s*\r?\n','')`\" -Verb RunAs
17 Exit
18 }
19
20 # Set console output encoding to UTF8 for clean symbols
21 [Console]::OutputEncoding = [System.Text.Encoding]::UTF8
22 Clear-Host
23
24 Write-Host "" -ForegroundColor Green
25 Write-Host " nullexit Windows Toggler " -ForegroundColor Green
26 Write-Host "" -ForegroundColor Green
27 Write-Host ""
28
29 # 2. Helpers
30 function Reset-HostDNS {
31 Write-Host "Restoring default DNS (DHCP/DHCP-assigned DNS)..." -ForegroundColor Cyan
32 $ActiveAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
33 foreach ($adapter in $ActiveAdapters) {
34 Write-Host " -> Resetting DNS on adapter: $($adapter.Name)" -ForegroundColor Gray
35 Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ResetServerAddresses -ErrorAction
36 SilentlyContinue
37 }
38 }
39
40 function Hijack-HostDNS ($ip) {
41 Write-Host "Hijacking DNS to route through AdGuard ($ip)..." -ForegroundColor Yellow
42 $ActiveAdapters = Get-NetAdapter | Where-Object { $_.Status -eq 'Up' }
43 foreach ($adapter in $ActiveAdapters) {
44 Write-Host " -> Pointing DNS to $ip on adapter: $($adapter.Name)" -ForegroundColor Gray
45 Set-DnsClientServerAddress -InterfaceIndex $adapter.InterfaceIndex -ServerAddresses $ip -ErrorAction
46 SilentlyContinue
47 }
48 }
49
50 function Get-HostTailscalePath {
51 $paths = @(
52 " $env:ProgramFiles\Tailscale\tailscale.exe",
53 " $env:ProgramFiles (x86)\Tailscale\tailscale.exe",
54 " tailscale.exe"
55 )
56 foreach ($path in $paths) {
57 if (Test-Path $path) { return $path }
58 }
59 return $null
60 }
```

```

57 # 3. Detect State
58 # Prevent deadlocks by resetting DNS to default temporarily during checks
59 Reset-HostDNS
60
61 Write-Host "Checking Docker daemon status..." -ForegroundColor Cyan
62 & docker info >$null 2>&1
63 if ($LASTEXITCODE -ne 0) {
64     Write-Host "Docker is not running. Starting Docker Desktop..." -ForegroundColor Yellow
65     $dockerPath = "C:\Program Files\Docker\Docker\Docker Desktop.exe"
66     if (Test-Path $dockerPath) {
67         Start-Process $dockerPath
68         Write-Host "Waiting for Docker daemon to become responsive..." -ForegroundColor Gray
69         $timeout = 60
70         while ($timeout -gt 0) {
71             Start-Sleep -Seconds 2
72             & docker info >$null 2>&1
73             if ($LASTEXITCODE -eq 0) {
74                 Write-Host "Docker is ready!" -ForegroundColor Green
75                 break
76             }
77             $timeout -= 2
78         }
79         if ($timeout -le 0) {
80             Write-Error "Docker failed to start in time. Aborting."
81             Exit 1
82         }
83     } else {
84         Write-Error "Docker Desktop executable not found at standard path. Please start Docker manually."
85         Exit 1
86     }
87 }
88
89 # Detect if gateway is already running
90 $runningWarp = docker compose ps --status running --format json | ConvertFrom-Json -ErrorAction
    SilentlyContinue | Where-Object { $_.Service -eq "warp" }
91 $isGatewayActive = ($null -ne $runningWarp)
92
93 # 4. Execution
94 if ($isGatewayActive) {
95     # STOP PATH
96     Write-Host "Gateway is currently RUNNING. Disabling now..." -ForegroundColor Yellow
97     Write-Host ""
98
99     # Disconnect host Tailscale from exit node if installed
100     $tsCli = Get-HostTailscalePath
101     if ($null -ne $tsCli) {
102         Write-Host "Disconnecting host Tailscale from exit node..." -ForegroundColor Cyan
103         & $tsCli up --exit-node= --accept-dns=true >$null 2>&1
104     }
105
106     Write-Host "Stopping Docker containers..." -ForegroundColor Cyan
107     docker compose down -t 5
108
109     Reset-HostDNS
110     Write-Host "Gateway has been successfully STOPPED." -ForegroundColor Green
111 } else {
112     # START PATH
113     Write-Host "Gateway is currently STOPPED. Enabling now..." -ForegroundColor Yellow
114     Write-Host ""
115
116     # Clean up corrupted AdGuardHome configurations
117     $confFile = "adguard\conf\AdGuardHome.yaml"
118     if (Test-Path $confFile) {
119         $fileSize = (Get-Item $confFile).Length
120         if ($fileSize -eq 0) {
121             Remove-Item $confFile -Force
122             Write-Host "Removed corrupted empty AdGuardHome.yaml to prevent container crash loop." -
                ForegroundColor Yellow
123         }
124     }
125
126     Write-Host "Starting Docker containers..." -ForegroundColor Cyan
127     docker compose up -d
128
129     Write-Host "Waiting for container's Tailscale connection to offer Exit Node..." -ForegroundColor Cyan
130     Write-Host " (This can take up to 60 seconds..." -ForegroundColor Gray
131
132     $elapsed = 0
133     $maxWait = 60
134     $resolvedIP = $null
135

```

```

136 while ($elapsed -lt $maxWait) {
137     $statusJson = docker compose exec -T tailscale tailscale status --json 2>$null
138     if ($null -ne $statusJson) {
139         $status = $statusJson | ConvertFrom-Json -ErrorAction SilentlyContinue
140         if ($null -ne $status -and $null -ne $status.Self) {
141             # Check if it has a valid Tailscale IP and is running
142             if ($status.Self.Online -and $status.Self.TailscaleIPs.Count -gt 0) {
143                 $resolvedIP = $status.Self.TailscaleIPs[0]
144                 break
145             }
146         }
147     }
148     Start-Sleep -Seconds 2
149     $elapsed += 2
150 }
151
152 if ($null -ne $resolvedIP) {
153     Write-Host "Tailscale IP resolved: $resolvedIP" -ForegroundColor Green
154     Write-Host ""
155
156     # Configure host Tailscale client to route through this exit node
157     $tsCli = Get-HostTailscalePath
158     if ($null -ne $tsCli) {
159         Write-Host "Configuring host Tailscale to use Exit Node ($resolvedIP)..." -ForegroundColor Cyan
160         & $tsCli up --exit-node=$resolvedIP --accept-dns=false >$null 2>&1
161     }
162
163     # Hijack DNS
164     Hijack-HostDNS $resolvedIP
165     Write-Host ""
166     Write-Host "Gateway has been successfully STARTED." -ForegroundColor Green
167 } else {
168     Write-Host "Warning: Could not resolve gateway Tailscale IP. DNS hijacking skipped." -ForegroundColor
169     Red
170     # Fallback to .gateway_ip if it exists
171     if (Test-Path ".gateway_ip") {
172         $fallbackIP = Get-Content ".gateway_ip" -TotalCount 1
173         if (-not [string]::IsNullOrEmpty($fallbackIP)) {
174             Write-Host "[Fallback] Using static IP from .gateway_ip: $fallbackIP" -ForegroundColor Yellow
175             Hijack-HostDNS $fallbackIP
176             Write-Host "Gateway STARTED (with static fallback IP)." -ForegroundColor Green
177         } else {
178             Write-Host "Gateway started in offline/unauthenticated state. Run setup.sh if this is a fresh setup
179             ." -ForegroundColor Yellow
180         }
181     }
182 }
183 Write-Host ""
184 Write-Host "Press any key to exit..."
185 [void]$Host.UI.RawUI.ReadKey("NoEcho,IncludeKeyDown")

```

## 31 scripts/check\_circuits.py

```

1  #!/usr/bin/env python3
2  import socket
3  import os
4  import re
5
6  def query_tor(port, command, password=""):
7      try:
8          s = socket.socket()
9          s.settimeout(2)
10         s.connect(('127.0.0.1', port))
11         s.sendall(f"AUTHENTICATE \"{password}\"\\r\\n{command}\\r\\nQUIT\\r\\n".encode())
12         response = b""
13         while True:
14             chunk = s.recv(4096)
15             if not chunk:
16                 break
17             response += chunk
18         return response.decode()
19     except Exception as e:
20         return f"Error connecting to Tor Control Port: {e}"
21

```

```

22 def get_ports():
23     ports = {}
24     ports_file = "/tmp/nullexit-ports.env"
25     if os.path.exists(ports_file):
26         with open(ports_file, 'r') as f:
27             for line in f:
28                 if line.startswith("export "):
29                     # Split at first '=' to handle potential multiple '=' in value
30                     parts = line.replace("export ", "").strip().split("=", 1)
31                     if len(parts) == 2:
32                         key, val = parts
33                         try:
34                             ports[key] = int(val)
35                         except ValueError:
36                             ports[key] = val
37
38     return ports
39
40 def get_password(ports):
41     # 1. Try ports.env
42     password = ports.get("TOR_PASSWORD")
43     if password:
44         return password
45
46     # 2. Try .env
47     if os.path.exists(".env"):
48         try:
49             with open(".env", "r") as f:
50                 for line in f:
51                     if line.startswith("ADGUARD_PASSWORD="):
52                         return line.split("=", 1)[1].strip().strip('"\'')
53                     if line.startswith("TOR_PASSWORD="):
54                         return line.split("=", 1)[1].strip().strip('"\'')
55         except Exception:
56             pass
57
58     # 3. Fallback
59     return "nullexit"
60
61 def main():
62     ports = get_ports()
63     control_port = ports.get("TOR_CONTROL_PORT", 9051)
64     password = get_password(ports)
65
66     print(f"Connecting to Tor Control Port on 127.0.0.1:{control_port}...")
67     status = query_tor(control_port, "GETINFO circuit-status", password)
68
69     if "Error" in status:
70         print(status)
71         return
72
73     print("\nActive Tor Circuits:")
74     circuits = re.findall(r"(\d+)\s+BUILT\s+(.+?)\s+PURPOSE", status)
75     if not circuits:
76         print("No active built circuits found yet. Tor might still be bootstrapping.")
77         return
78
79     for circ_id, path_str in circuits:
80         print(f"\nCircuit #{circ_id}:")
81         nodes = path_str.split(",")
82         for idx, node in enumerate(nodes):
83             fingerprint = node.split("-")[0].replace("$", "")
84             nickname = node.split("-")[1] if "~" in node else "unknown"
85
86             # Fetch node IP
87             ns_info = query_tor(control_port, f"GETINFO ns/id/{fingerprint}", password)
88             ip = "unknown"
89             for line in ns_info.split("\n"):
90                 if line.startswith("r "):
91                     parts = line.split(" ")
92                     if len(parts) >= 7:
93                         ip = parts[6]
94                         break
95
96             # Fetch country
97             country = "unknown"
98             if ip != "unknown":
99                 country_info = query_tor(control_port, f"GETINFO ip-to-country/{ip}", password)
100                 for line in country_info.split("\n"):
101                     if line.startswith("250-ip-to-country/"):
102                         country = line.split("=")[1].strip().upper()
103                         break

```



```

103         role = "Guard" if idx == 0 else ("Exit" if idx == len(nodes) - 1 else "Middle")
104         print(f" [{role}] {nickname} ({ip}) - Country: {country}")
105
106
107 if __name__ == "__main__":
108     main()

```

## 32 scripts/diagnose-host-leak.sh

```

1 #!/bin/bash
2 # scripts/diagnose-host-leak.sh nullexit host-routing diagnostic
3 #
4 # Symptom this addresses: tests like https://www.whatismyip.com on the
5 # *HOST MAC* return the underlying physical ISP/ASN
6 # local ISP / campus network") even though the nullexit gateway is active and remote
7 # tailnet clients correctly egress via Cloudflare WARP. The container and
8 # the gateway itself are healthy only the host's own routing is leaking.
9 #
10 # This script runs 7 checks in sequence, classifies the result into ONE
11 # of the three known leak scenarios, writes the full report to a
12 # timestamped file in the project root, and prints a verdict + a
13 # ready-to-run fix command on stdout. Pass --fix to apply the matching
14 # remediation in the same invocation and re-verify egress afterwards.
15 # See devref.md 10.30 for the deep-dive rationale.
16 #
17 # Usage:
18 #   bash scripts/diagnose-host-leak.sh           # diagnose + write report
19 #   bash scripts/diagnose-host-leak.sh --fix      # diagnose + fix + re-verify
20 #   bash scripts/diagnose-host-leak.sh --watch    # full baseline then loop checks every 60s
21 #   bash scripts/diagnose-host-leak.sh --watch 30 # same, with custom interval (seconds)
22 #   bash scripts/diagnose-host-leak.sh --help    # this message
23 #
24 # Multiple runs produce timestamped files (one per run) so historical
25 # reports can be diffed. The newest file is the most recent.
26
27 set -uo pipefail
28 # NOTE: errexit (`set -e`) is intentionally NOT enabled. Several checks
29 # below are EXPECTED to fail and that's diagnostic data (e.g. `tailscale
30 # status` returns non-zero when the daemon is wedged). Every command
31 # pipes to a helper that records the exit code without killing the script.
32
33 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
34 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
35 TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
36 source "$SCRIPT_DIR/common.sh"
37 LOG_FILE="$PROJECT_ROOT/output.log"
38 REPORT_FILE="$PROJECT_ROOT/host-leak-diagnostic-$TIMESTAMP.txt"
39
40 # Argument parsing
41 DO_FIX=false
42 DO_WATCH=false
43 WATCH_INTERVAL=60
44 WATCH_LOG="$PROJECT_ROOT/host-leak-watch.log"
45 case "${1:-}" in
46     --fix) DO_FIX=true ;;
47     --watch)
48         DO_WATCH=true
49         # Optional second argument: custom interval in seconds
50         if [ -n "${2:-}" ] && [[ "$2" =~ ^[0-9]+$ ]]; then
51             WATCH_INTERVAL="$2"
52         fi
53     ;;
54     --help|-h)
55         sed -n '2,31p' "$0"
56         exit 0
57     ;;
58     *)
59         : # default: diagnose only
60     ;;
61 *)
62     echo "Unknown argument: $1" >&2
63     echo "Run with --help for usage." >&2
64     exit 2
65 ;;
66 esac
67

```

```

68 # Colours (auto-disabled when stdout is not a tty, e.g. piped)
69 if [ -t 1 ]; then
70     RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
71     BOLD='\033[1m'; NC='\033[0m'
72 else
73     RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
74 fi
75
76 # Output helpers write to BOTH stdout and the report file
77 exec 3>> "$REPORT_FILE" || {
78     printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
79     exit 1
80 }
81 section() {
82     local title="$1"
83     printf '\n\b %s %b\n' "$BOLD" "$title" "$NC" >&3
84     printf '\n\b %s %b\n' "$BOLD" "$title" "$NC"
85 }
86 line() { printf ' %b\n' "$1" >&3; printf ' %b\n' "$1"; }
87 ok() { line "${GREEN}${NC} $1"; }
88 warn() { line "${YELLOW}${NC} $1"; }
89 fail() { line "${RED}${NC} $1"; }
90 note() { line "(no ${1})"; }
91
92 # Watch-mode header (only shown when entering watch loop)
93 if [ "$DO_WATCH" = "true" ]; then
94     echo ""
95     echo ""
96     echo " n u l l e x i t   w a t c h   m o d e "
97     echo ""
98     echo " Interval:          ${WATCH_INTERVAL}s"
99     echo " Watch log:          $WATCH_LOG"
100     echo ""
101 else
102     echo ""
103     echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c "
104     echo ""
105     echo " Timestamp (UTC): $TIMESTAMP"
106     echo " Project root:    $PROJECT_ROOT"
107     echo " Report file:     $REPORT_FILE"
108     echo " Mode:            ${([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only (pass --
109         fix to remediate)")}
109     echo ""
110     echo "" >&3
111     echo "" >&3
112     echo " n u l l e x i t   h o s t - r o u t i n g   d i a g n o s t i c " >&3
113     echo "" >&3
114     echo " Timestamp (UTC): $TIMESTAMP" >&3
115     echo " Mode:            ${([ "$DO_FIX" = true ] && echo "diagnose + apply fix" || echo "diagnose only")} >&3
116 fi
117
118 # Always add Homebrew to PATH (same as toggle.sh / recover.sh)
119 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
120
121 # Pure-bash timeout (macOS lacks GNU `timeout`)
122 run_with_timeout() {
123     local timeout_sec="$1"
124     shift
125     if [ "$#" -eq 0 ]; then return 1; fi
126     "$@" &
127     local cmd_pid=$!
128     (
129         sleep "$timeout_sec"
130         if kill -0 "$cmd_pid" 2>> "$LOG_FILE"; then
131             kill -9 "$cmd_pid" 2>> "$LOG_FILE" || true
132         fi
133     ) >> "$LOG_FILE" 2>&1 &
134     local watcher_pid=$!
135     set +e
136     wait "$cmd_pid"
137     local exit_status=$?
138     set -uo pipefail
139     kill "$watcher_pid" 2>> "$LOG_FILE" || true
140     return "$exit_status"
141 }
142
143 #
144 # Quick-check functions (used by --watch loop; return simple status strings)
145 # Each echoes its result so the caller can capture and compare.
146 #
147

```

```

148 quick_warp() {
149     # Returns: "on", "off", "unknown"
150     local out
151     out=$(run_with_timeout 8 curl -s --max-time 6 \
152         https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null || echo "")
153     local warp
154     warp=$(printf '%s' "$out" | awk -F=' ' '/^warp={print $2; exit}')
155     printf '%s' "${warp:-unknown}"
156 }
157
158 quick_ipv6() {
159     # Returns: the IPv6 address if leaking, or empty string if clean
160     local out
161     out=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>/dev/null || echo "")
162     if [ -n "$out" ] && printf '%s' "$out" | grep -qE '[0-9a-fA-F:]+$'; then
163         printf '%s' "$out"
164     fi
165 }
166
167 quick_default_route() {
168     # Returns: "utun" if default route is via Tailscale utun*,
169     #           "wifi" if via physical Wi-Fi (192.168.x.x),
170     #           "other" otherwise
171     local route
172     route=$(netstat -rn 2>/dev/null | awk '/^default/ {print $0; exit}')
173     if [ -z "$route" ]; then
174         printf '%s' "none"
175     elif printf '%s' "$route" | grep -qE 'utun[0-9]+'; then
176         printf '%s' "utun"
177     elif printf '%s' "$route" | grep -qE '^(default|0\.\0\.\0\.\0).*\b192\.\168\.'; then
178         printf '%s' "wifi"
179     else
180         printf '%s' "other"
181     fi
182 }
183
184 # Watch loop (runs after baseline diagnostic when --watch is passed)
185 run_watch_loop() {
186     local baseline_warp="$1"
187     local baseline_default="$2"
188     local baseline_ipv6_leak="$3"
189     local interval="$4"
190
191     echo ""
192     echo ""
193     echo " Baseline established. Monitoring every ${interval}s..."
194     echo "   warp:      $baseline_warp"
195     echo " default:    $baseline_default"
196     echo ""
197
198     # Safety: warn if interval is very short (avoids hammering APIs)
199     if [ "$interval" -lt 30 ]; then
200         printf '%b' "Short interval (%ss) rate-limiting risk on external APIs%b\n" "$YELLOW" "$interval" "$NC"
201         printf ' Consider 30s or longer for production use.\n'
202     fi
203     printf '[%s] WATCH START warp=%s default=%s interval=%ss\n' \
204         "$(date -u +%Y-%m-%dT%H:%M:%SZ)" "$baseline_warp" "$baseline_default" "$interval" \
205         >> "$WATCH_LOG"
206
207     local iteration=0
208     local last_warp="$baseline_warp"
209     local last_default="$baseline_default"
210     local last_ipv6_ok=$( [ "$baseline_ipv6_leak" = "false" ] && echo true || echo false )
211
212     trap 'printf "\n%bWatch stopped by user. Log at: %s%b\n" "$YELLOW" "$WATCH_LOG" "$NC"; exit 0' INT TERM
213
214     while true; do
215         iteration=$((iteration + 1))
216
217         local warp_now ipv6_now default_now
218         warp_now=$(quick_warp)
219         ipv6_now=$(quick_ipv6)
220         default_now=$(quick_default_route)
221
222         local ts
223         ts=$(date -u +%H:%M:%S)
224
225         # WARP check
226         if [ "$warp_now" != "$last_warp" ]; then
227             if [ "$warp_now" = "off" ] || [ "$warp_now" = "unknown" ]; then
228                 printf '\n%b ALERT %b\n' "$RED$BOLD" "$NC"

```

```

229     printf '%b[%s] %bWARP FLIPPED: %s %s\b YOUR IP IS LEAKING!\n' \
230         "$RED" "$ts" "$BOLD" "$last_warp" "$warp_now" "$NC"
231     printf '[%s] ALERT warp flipped %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
232 else
233     printf '\n\b[%s] warp recovered: %s %s\b\n' \
234         "$GREEN" "$ts" "$last_warp" "$warp_now" "$NC"
235     printf '[%s] OK warp recovered %s%s\n' "$ts" "$last_warp" "$warp_now" >> "$WATCH_LOG"
236 fi
237 last_warp="$warp_now"
238 fi
239
240 # IPv6 check
241 if [ -n "$ipv6_now" ]; then
242     if [ "$last_ipv6_ok" = true ]; then
243         printf '\n\b ALERT %b\n' "$RED$BOLD" "$NC"
244         printf '%b[%s] %bIPv6 LEAK DETECTED %s\b\n' \
245             "$RED" "$ts" "$BOLD" "$ipv6_now" "$NC"
246         printf '[%s] ALERT IPv6 leak %s\n' "$ts" "$ipv6_now" >> "$WATCH_LOG"
247         last_ipv6_ok=false
248     fi
249 else
250     if [ "$last_ipv6_ok" = false ]; then
251         printf '%b[%s] IPv6 leak resolved\b\n' "$GREEN" "$ts" "$NC"
252         printf '[%s] OK IPv6 resolved\n' "$ts" >> "$WATCH_LOG"
253     fi
254     last_ipv6_ok=true
255 fi
256
257 # Default route check
258 if [ "$default_now" != "$last_default" ]; then
259     if [ "$default_now" != "utun" ]; then
260         printf '\n\b ALERT %b\n' "$RED$BOLD" "$NC"
261         printf '%b[%s] %bDEFAULT ROUTE CHANGED: %s %s\b traffic may bypass WARP!\n' \
262             "$RED" "$ts" "$BOLD" "$last_default" "$default_now" "$NC"
263         printf '[%s] ALERT route changed %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
264     else
265         printf '%b[%s] default route recovered: %s %s\b\n' \
266             "$GREEN" "$ts" "$last_default" "$default_now" "$NC"
267         printf '[%s] OK route recovered %s%s\n' "$ts" "$last_default" "$default_now" >> "$WATCH_LOG"
268     fi
269     last_default="$default_now"
270 fi
271
272 # Heartbeat (every 10th iteration or if all OK)
273 if [ $((iteration % 10)) -eq 0 ]; then
274     printf '[%s] #d warp=%s ipv6=%s route=%s\n' \
275         "$ts" "$iteration" "$warp_now" "${ipv6_now:-none}" "$default_now" >> "$WATCH_LOG"
276     printf '\r\b[%s] #d warp=%s route=%s (Ctrl+C to stop)%b\n' \
277         "$GREEN" "$ts" "$iteration" "$warp_now" "$default_now" "$NC"
278 fi
279
280 sleep "$interval"
281 done
282 }
283
284 #
285 # Begin main diagnostic (skipped in non-watch modes if we're just watching)
286 #
287
288 ACTIVE_SERVICE=$(get_active_service)
289 ENO_SERVICE=$(networksetup -listnetworkserviceorder 2>> "$LOG_FILE" \
290     | grep -B 1 "Device: en0" | head -1 \
291     | sed -E 's/^\([0-9\+]+\)/ /' || true)
292 [ -z "$ENO_SERVICE" ] && ENO_SERVICE="Wi-Fi"
293
294 # Resolve the gateway's Tailscale IP.
295 GATEWAY_TS_IP=""
296 if [ -f "$PROJECT_ROOT/.gateway_ip" ]; then
297     GATEWAY_TS_IP=$(cat "$PROJECT_ROOT/.gateway_ip" 2>> "$LOG_FILE" \
298         | tr -d '\r' | awk 'NR==1{print $1; exit}')
299 fi
300 if [ -z "$GATEWAY_TS_IP" ] && command -v docker >> "$LOG_FILE" 2>&1 \
301     && docker compose ps --status running 2>/dev/null | grep -q 'tailscale'; then
302     GATEWAY_TS_IP=$(run_with_timeout 10 docker compose exec -T tailscale \
303         tailscale ip -4 2>> "$LOG_FILE" \
304         | tr -d '\r' | awk 'NR==1{print $1; exit}')
305 fi
306
307 # Scenario classification state
308 SCENARIO=""
309 WARP_STATUS=""

```

```

310 SOCKS5_ENABLED=false
311 TS_ON_MESH=false
312 DEFAULT_VIA_UTUN=false
313 IPV6_LEAK=false
314
315 #
316 section "1/8 Host Tailscale mesh status"
317 #
318 if ! command -v tailscale >> "$LOG_FILE" 2>&1; then
319     fail "tailscale CLI not on PATH install via 'brew install tailscale' then re-run"
320     echo " (a Tailscale daemon is required for the exit-node path to work)"
321 else
322     set +e
323     TS_OUT=$(run_with_timeout 5 tailscale status 2>&1)
324     TS_EXIT=$?
325     set -uo pipefail
326     TS_HEAD=$(printf '%s\n' "$TS_OUT" | head -5 | tr -d '\r')
327     if [ "$TS_EXIT" -eq 137 ]; then
328         fail "tailscaled daemon is wedged/unresponsive (5s timeout)"
329         echo " Fix path: 'sudo brew services restart tailscale'"
330     elif printf '%s' "$TS_OUT" | grep -qE 'Logged out|not logged in|expired'; then
331         fail "tailscale is logged out / auth expired on this host"
332         echo " Fix path: 'sudo tailscale up' (browser auth once)"
333     elif printf '%s' "$TS_OUT" | grep -q '100\.'; then
334         ok "Host is on tailnet (100.x.x.x visible in status)"
335         TS_ON_MESH=true
336         if printf '%s' "$TS_OUT" | grep -q "$GATEWAY_TS_IP"; then
337             ok "Gateway $GATEWAY_TS_IP visible in host's peer list"
338         elif [ -n "$GATEWAY_TS_IP" ]; then
339             warn "Gateway $GATEWAY_TS_IP NOT in host's peer list possible auth/key mismatch"
340         fi
341         line " first 5 lines of 'tailscale status' "
342         line "$(printf '%s' "$TS_HEAD" | awk '{print " " "$0}')"
343     else
344         fail "tailscale output unrecognized:"
345         line "$(printf '%s' "$TS_OUT" | head -3 | awk '{print " " "$0}')"
346     fi
347 fi
348
349 #
350 section "2/8 Active SOCKS5 proxy on Wi-Fi"
351 #
352 SOCKS_OUT=$(networksetup -getsocksfirewallproxy "$ACTIVE_SERVICE" 2>> "$LOG_FILE" \
353     || networksetup -getsocksfirewallproxy Wi-Fi 2>> "$LOG_FILE" \
354     || echo "Could not query SOCKS5 proxy state")
355 if printf '%s' "$SOCKS_OUT" | grep -q "Enabled: Yes"; then
356     fail "SOCKS5 proxy IS enabled on $ACTIVE_SERVICE toggle.sh fell back to SOCKS5 mode last run"
357     SOCKS5_ENABLED=true
358     line " Full state:"
359     line "$(printf '%s' "$SOCKS_OUT" | awk '{print " " "$0}')"
360     line " Browse on macOS ignores system SOCKS5 by default traffic bypasses WARP."
361 else
362     ok "SOCKS5 proxy is OFF on $ACTIVE_SERVICE exit-node path should be active"
363 fi
364
365 #
366 section "3/8 Egress IP + WARP tunnel verification (definitive test)"
367 #
368 CDN_OUT=$(run_with_timeout 8 curl -s --max-time 6 \
369     https://www.cloudflare.com/cdn-cgi/trace 2>&1 || echo "(curl failed)")
370 CDN_WARP=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^warp={print $2; exit}')
371 CDN_IP=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^ip={print $2; exit}')
372 CDN_COLO=$(printf '%s' "$CDN_OUT" | awk -F=' ' '/^colo={print $2; exit}')
373 WARP_STATUS="$CDN_WARP"
374 if [ "$CDN_WARP" = "on" ]; then
375     ok "warp=on tunnel is hot, host traffic IS going through Cloudflare WARP"
376     line " Public IP: $CDN_IP"
377     line " Cloudflare: ${CDN_COLO:-unknown colo}"
378     line " whatismyip.com's 'local ISP' result is its ISP-database error, not yours."
379 elif [ "$CDN_WARP" = "off" ]; then
380     fail "warp=off host traffic is NOT going through WARP"
381     line " Public IP: $CDN_IP (NOT Cloudflare)"
382     warn "This is the actual leak."
383 else
384     warn "could not determine warp status (curl failed or returned unexpected shape)"
385     line " Raw output: $(printf '%s' "$CDN_OUT" | head -3 | tr '\n' '|')"
386 fi
387
388 #
389 section "4/8 IPv6 leak probe (bypasses IPv4 exit-node on campus networks)"
390 #

```

```

391 IPV6_OUT=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
392 if [ -n "$IPV6_OUT" ] && printf '%s' "$IPV6_OUT" | grep -qE '[0-9a-fA-F:]+$'; then
393     fail "host has live IPv6 egress $IPV6_OUT (A6 record bypasses WARP)"
394     IPV6_LEAK=true
395     line " Tailscale exit-node advertisement is IPv4-only, so IPv6 traffic skips it."
396     line " Quick-fix: 'sudo networksetup -setv6off $ACTIVE_SERVICE'"
397 else
398     ok "no IPv6 egress detected (good IPv6 won't bypass the tunnel)"
399 fi
400
401 #
402 section "5/8 Host egress route for real traffic (route -n get 1.1.1.1)"
403 #
404 # WHY NOT 'netstat -rn | grep default':
405 # On macOS, Tailscale does NOT replace the physical default route installed by
406 # DHCP. Instead it adds a second interface-scoped route bound to the utun
407 # interface that the kernel prefers for real traffic forwarding. The literal
408 # "default" entry (192.168.x.x via en0) is left untouched in the table as a
409 # BSD routing quirk it's not lying, it's just answering a different question.
410 # The correct question is: "where does a real destination actually resolve?"
411 # `route -n get 1.1.1.1` asks the kernel's forwarding logic directly.
412 ROUTE_GET=$(route -n get 1.1.1.1 2>/dev/null)
413 ROUTE_IFACE=$(echo "$ROUTE_GET" | awk '/interface:/{print $2}')
414 ROUTE_GATEWAY=$(echo "$ROUTE_GET" | awk '/gateway:/{print $2}')
415 if [ -z "$ROUTE_IFACE" ]; then
416     warn "route -n get 1.1.1.1 returned no interface routing table may be empty"
417 else
418     line " interface: $ROUTE_IFACE gateway: ${ROUTE_GATEWAY:-n/a}"
419     if echo "$ROUTE_IFACE" | grep -qE '^utun[0-9]+'; then
420         ok "1.1.1.1 resolves via $ROUTE_IFACE Tailscale interface-scoped route is winning"
421         DEFAULT_VIA_UTUN=true
422     elif echo "$ROUTE_IFACE" | grep -qE '^en[0-9]+'; then
423         fail "1.1.1.1 resolves via $ROUTE_IFACE (physical Wi-Fi) exit-node interface route not active"
424     else
425         warn "1.1.1.1 resolves via unexpected interface: $ROUTE_IFACE"
426     fi
427 fi
428
429 #
430 section "6/8 Last toggle.sh START-relevant lines in output.log (if any)"
431 #
432 if [ ! -f "$LOG_FILE" ]; then
433     note "output.log"
434     line " toggle.sh was never run from this directory. Run './toggle.sh' first,"
435     line " then re-run this diagnostic."
436 else
437     HITS=$(grep -E 'Exit node enabled|SKIP_EXIT_NODE|Pre-flight|Falling back to SOCKS|Starting|Successfully' \
438 "$LOG_FILE" 2>/dev/null | tail -8 || true)
439     if [ -z "$HITS" ]; then
440         note "matching lines in output.log"
441     else
442         line "$(printf '%s\n' "$HITS" | awk '{print " " "$0}')"
443     fi
444 fi
445
446 #
447 section "7/8 Gateway IP we should be pointing at"
448 #
449 if [ -n "$GATEWAY_TS_IP" ]; then
450     ok "gateway Tailscale IP: $GATEWAY_TS_IP"
451 else
452     warn "could not resolve gateway Tailscale IP"
453     line " (looked in .gateway_ip and tried docker compose exec tailscale ip -4)"
454 fi
455
456 #
457 section "8/8 Tailscale System Extension (App Store conflict check)"
458 #
459 SE_PENDING_UNINSTALL=false
460 SE_ACTIVE=false
461 if command -v systemextensionsctl >/dev/null 2>&1; then
462     SE_OUT=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale')
463     if [ -n "$SE_OUT" ]; then
464         line " Tailscale System Extension(s) detected:"
465         line "$(printf '%s' "$SE_OUT" | awk '{print " " "$0}')"
466         if printf '%s' "$SE_OUT" | grep -q 'waiting to uninstall'; then
467             SE_PENDING_UNINSTALL=true
468             warn "System Extension is pending uninstall on next reboot."
469         else
470             SE_ACTIVE=true
471             warn "Active Tailscale System Extension detected. This will conflict with Homebrew Tailscale."

```

```

472     fi
473   else
474     ok "no Tailscale System Extension detected"
475   fi
476 else
477   note "systemextensionsctl not available (not on macOS or permission restricted)"
478 fi
479
480 #
481 section "CLASSIFICATION    applies ONE of the four known scenarios"
482 #
483
484 if [ "$WARP_STATUS" = "on" ] && [ "$IPV6_LEAK" = "false" ]; then
485   SCENARIO="OK"
486 elif [ "$SE_PENDING_UNINSTALL" = "true" ]; then
487   SCENARIO="SE_PENDING"
488 elif [ "$SOCKS5_ENABLED" = "true" ]; then
489   SCENARIO="A"
490 elif [ "$IPV6_LEAK" = "true" ]; then
491   SCENARIO="B"
492 elif [ "$DEFAULT_VIA_UTUN" = "false" ] && [ "$TS_ON_MESH" = "true" ]; then
493   SCENARIO="C"
494 else
495   SCENARIO="UNKNOWN"
496 fi
497
498 case "$SCENARIO" in
499   SE_PENDING)
500     echo ""
501     echo -e " ${RED}${BOLD}VERDICT: Scenario SE_PENDING Tailscale System Extension Pending Uninstall${NC}"
502     echo " A prior App Store Tailscale System Extension is pending uninstall."
503     echo " Since System Integrity Protection (SIP) is enabled, macOS blocks manual"
504     echo " uninstallation of terminated system extensions until the next reboot."
505     echo " Brew-managed tailscaled cannot engage the Network Extension slot until this is resolved."
506     echo ""
507     echo " ${BOLD}Recommended fix:${NC}"
508     echo "     sudo shutdown -r now"
509     echo "     # After reboot, run: ./toggle.sh"
510     ;;
511   A)
512     echo ""
513     echo -e " ${RED}${BOLD}VERDICT: Scenario A SOCKS5 fallback lane${NC}"
514     echo " toggle.sh's pre-flight checks failed last run. The script activated"
515     echo " SOCKS5 + local DNS proxy as a fallback. DNS gets hijacked but"
516     echo " browsers on macOS ignore the system SOCKS5 traffic egresses"
517     echo " direct through the local network."
518     echo ""
519     echo " ${BOLD}Recommended fix:${NC}"
520     echo "     echo 'GATEWAY_BYPASS_PING=true' >> $PROJECT_ROOT/.env"
521     echo "     sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \\"
522     if [ -n "$GATEWAY_TS_IP" ]; then
523       echo "         --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
524     else
525       echo "         --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true # fill in $GATEWAY_TS_IP"
526     fi
527     echo "     sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
528     echo "     ./toggle.sh"
529     ;;
530   B)
531     echo ""
532     echo -e " ${RED}${BOLD}VERDICT: Scenario B IPv6 leak over dual-stack campus/local IPv6${NC}"
533     echo " Tailscale exit-node advertises only IPv4. macOS sends AAAA queries"
534     echo " straight out en0, which is dual-stack on most campus APs IPv6"
535     echo " traffic bypasses the WARP tunnel entirely."
536     echo ""
537     echo " ${BOLD}Recommended fix:${NC}"
538     echo "     sudo networksetup -setv6off $ACTIVE_SERVICE"
539     echo "     sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
540     echo "     # Re-enable IPv6 later with: sudo networksetup -setv6automatic $ACTIVE_SERVICE"
541     ;;
542   C)
543     echo ""
544     echo -e " ${RED}${BOLD}VERDICT: Scenario C Tailscale route-freeze${NC}"
545     echo " Host's default route points at the Wi-Fi gateway instead of the"
546     echo " Tailscale utun* interface. The exit-node preference is set but the"
547     echo " routing-table assertion did not take effect (devref 10.26)."
548     echo ""
549     echo " ${BOLD}Recommended fix:${NC}"
550     echo "     sudo brew services restart tailscale"
551     echo "     sudo route delete -net 0.0.0.0/1"
552     echo "     sudo route delete -net 128.0.0.0/1"

```



```

553     echo "        sudo dscacheutil -flushcache && sudo killall -HUP mDNSResponder"
554     if [ -n "$GATEWAY_TS_IP" ]; then
555     echo "        sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \"
556     echo "            --exit-node=$GATEWAY_TS_IP --exit-node-allow-lan-access=true"
557     else
558     echo "        sudo tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true --exit-node=$TS_IP
559     echo "            --exit-node-allow-lan-access=true"
560     fi
561     ;;
562 OK)
563     echo ""
564     echo -e "    ${GREEN}${BOLD}VERDICT: No host leak detected tunnel is working${NC}"
565     echo "    The local ISP string from whatismyip.com is its IP-database"
566     echo "    misclassifying the Cloudflare egress IP. CDN-CGI's trace endpoint"
567     echo "    (which is what we use here) reports warp=on because that endpoint"
568     echo "    is hosted by Cloudflare itself and can see its own edge."
569     echo "    If you want a third-party confirmation:"
570     echo "    curl -s https://api.ipify.org; echo"
571     ;;
572 *)
573     echo ""
574     echo -e "    ${YELLOW}${BOLD}VERDICT: Could not classify automatically${NC}"
575     echo "    The diagnostic data did not match any known scenario cleanly."
576     echo "    Likely causes:"
577     echo "        - tailscaled is logged out / expired (see Section 1)"
578     echo "        - Docker is not running (gateway containers down)"
579     echo "        - .gateway_ip missing AND docker compose exec failed"
580     echo "    Inspect the report at: $REPORT_FILE"
581     ;;
582 esac
583 echo ""
584
585 #
586 # Write the verdict block to the report file
587 #
588 SECOND_BLOCK=$(cat <<EOF
589
590
591 C L A S S I F I C A T I O N   R E S U L T
592
593 Scenario: $SCENARIO
594 WARP egress: $WARP_STATUS (cdn-cgi/trace)
595 SOCKS5 lane: $([ "$SOCKS5_ENABLED" = true ] && echo "ENABLED (browsers bypass)" || echo "off")
596 IPv6 leak: $([ "$IPV6_LEAK" = true ] && echo "YES (campus IPv6 egress)" || echo "no")
597 Default via: $([ "$DEFAULT_VIA_UTUN" = true ] && echo "utun* (Tailscale)" || echo "Wi-Fi gateway")
598 Host on mesh: $([ "$TS_ON_MESH" = true ] && echo "yes" || echo "no")
599 SE pending: $([ "$SE_PENDING_UNINSTALL" = true ] && echo "yes" || echo "no")
600 Gateway IP: ${GATEWAY_TS_IP:-unresolved}
601 EOF
602 )
603 printf '%s\n' "$SECOND_BLOCK" >&3
604
605 #
606 # --fix mode
607 #
608 if [ "$DO_FIX" = "true" ] && [ "$SCENARIO" != "OK" ] && [ "$SCENARIO" != "UNKNOWN" ]; then
609     echo ""
610     echo ""
611     echo "    A P P L Y I N G   F I X (scenario $SCENARIO)"
612     echo ""
613     echo ""
614
615     case "$SCENARIO" in
616         SE_PENDING)
617             warn "System Extension uninstall is pending on next reboot."
618             line "Please run 'sudo shutdown -r now' to reboot the system."
619             ;;
620         A)
621             if [ -n "$GATEWAY_TS_IP" ]; then
622                 line "writing GATEWAY_BYPASS_PING=true to .env"
623                 ENV="$PROJECT_ROOT/.env"
624                 [ -f "$ENV" ] || touch "$ENV"
625                 if grep -q '^GATEWAY_BYPASS_PING=' "$ENV" 2>/dev/null; then
626                     sed -i '' 's/^GATEWAY_BYPASS_PING=.*GATEWAY_BYPASS_PING=true/' "$ENV"
627                 else
628                     printf '\nGATEWAY_BYPASS_PING=true\n' >> "$ENV"
629                 fi
630                 line "re-asserting tailscale exit-node (with --reset)"
631                 sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
632                     --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \

```

```

633     >> "$LOG_FILE" 2>&1 \
634     && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
635     || warn "tailscale up did not return success (check $LOG_FILE)"
636 else
637     warn "no gateway Tailscale IP resolved skipping tailscale up"
638     warn "fix manually: 'sudo tailscale up --reset ... --exit-node=<TS_IP>'"
639 fi
640 ;;
641 B)
642     line "disabling IPv6 on $ACTIVE_SERVICE while gateway is up"
643     sudo -n networksetup -setv6off "$ACTIVE_SERVICE" >> "$LOG_FILE" 2>&1 \
644     && ok "IPv6 disabled on $ACTIVE_SERVICE (re-enable with 'setv6automatic')\" \
645     || warn "networksetup -setv6off failed (check $LOG_FILE)"
646 ;;
647 C)
648     line "restarting tailscaled (route-table fix)"
649     run_with_timeout 15 brew services restart tailscale >> "$LOG_FILE" 2>&1 \
650     || run_with_timeout 15 sudo -n brew services restart tailscale >> "$LOG_FILE" 2>&1
651     sleep 3
652     line "flushing stale host routes + DNS cache"
653     sudo -n route delete -net 0.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
654     sudo -n route delete -net 128.0.0.0/1 >> "$LOG_FILE" 2>&1 || true
655     sudo -n dscacheutil -flushcache >> "$LOG_FILE" 2>&1 || true
656     sudo -n killall -HUP mDNSResponder >> "$LOG_FILE" 2>&1 || true
657     if [ -n "$GATEWAY_TS_IP" ]; then
658         line "re-asserting exit-node after restart"
659         sudo -n tailscale up --reset --ssh=true --accept-dns=false --accept-routes=true \
660         --exit-node="$GATEWAY_TS_IP" --exit-node-allow-lan-access=true \
661         >> "$LOG_FILE" 2>&1 \
662         && ok "tailscale exit-node re-asserted for $GATEWAY_TS_IP" \
663         || warn "tailscale up did not return success (check $LOG_FILE)"
664     fi
665 ;;
666 esac
667
668 echo ""
669 echo "Re-verifying after fix..."
670 sleep 3
671 CDN_WARP_AFTER=$(run_with_timeout 8 curl -s --max-time 6 \
672     https://www.cloudflare.com/cdn-cgi/trace 2>&1 \
673     | awk -F=' ' /^warp=/{print $2; exit}')
674 IPV6_AFTER=$(run_with_timeout 7 curl -6 -s --max-time 5 ifconfig.co 2>&1 || echo "")
675 printf '\n[AFTER FIX]\n' >&3
676 printf ' warp status: %s\n' "$CDN_WARP_AFTER" >&3
677 printf ' IPv6 egress: %s\n' "${IPV6_AFTER:-none}" >&3
678 line " warp status: ${CDN_WARP_AFTER:-unknown}"
679 line " IPv6 egress: ${IPV6_AFTER:-none}"
680 if [ "$CDN_WARP_AFTER" = "on" ] && [ -z "$IPV6_AFTER" ]; then
681     echo ""
682     ok "Fix verified host is now routing through WARP. Remote clients unaffected."
683 else
684     echo ""
685     warn "Fix did not fully take effect on first try. Likely causes:"
686     line " tailscaled needed a longer settling window (re-run diagnostic in 30s)"
687     line " The classified scenario was wrong; paste this report for triage."
688     line " Tailscale needs a full auth refresh: 'sudo tailscale up'"
689 fi
690 fi
691
692 echo ""
693 echo ""
694 echo " Full report written to:"
695 echo " $REPORT_FILE"
696 echo ""
697 echo ""
698
699 #
700 # --watch mode: enter the continuous monitoring loop after baseline diagnostic
701 #
702 if [ "$DO_WATCH" = "true" ]; then
703     # Build baseline from the full diagnostic just completed
704     BASELINE_WARP="$WARP_STATUS"
705     BASELINE_DEFAULT=$( [ "$DEFAULT_VIA_UTUN" = "true" ] && echo "utun" || echo "wifi/other")
706
707     if [ "$SCENARIO" != "OK" ]; then
708         warn "Baseline diagnostic found issues (scenario: $SCENARIO)."
709         line " Watch loop will still run alerts fire on any STATE CHANGE from current baseline."
710         line " Fix the issue first? Re-run with: bash scripts/diagnose-host-leak.sh --fix"
711         echo ""
712     fi
713

```

```

714     run_watch_loop "$BASELINE_WARP" "$BASELINE_DEFAULT" "$IPV6_LEAK" "$WATCH_INTERVAL"
715 fi
716
717 exit 0

```

## 33 scripts/dns-proxy.py

```

1  #!/usr/bin/env python3
2  """
3  dns-proxy.py Local DNS proxy for nullexit gateway.
4
5  Listens on UDP port 53, receives DNS queries, forwards them over TCP
6  to AdGuard at 127.0.0.1:5354 (the Docker port mapping), and returns
7  the response via UDP. Handles the DNS-over-TCP wire format (2-byte
8  length prefix) correctly unlike socat which just forwards raw bytes.
9  """
10
11 import socket
12 import struct
13 import sys
14 import os
15 from concurrent.futures import ThreadPoolExecutor
16
17 LISTEN_ADDR = os.environ.get("LISTEN_ADDR", "127.0.0.1")
18 LISTEN_PORT = int(os.environ.get("LISTEN_PORT", 53))
19 TARGET_HOST = os.environ.get("TARGET_HOST", "127.0.0.1")
20 TARGET_PORT = int(os.environ.get("TARGET_PORT", 5354))
21
22
23 def handle_query(data: bytes, client_addr: tuple, sock: socket.socket) -> None:
24     """Forward a single DNS query over TCP and send the response back via UDP."""
25     try:
26         tcp = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
27         tcp.settimeout(5)
28         tcp.connect((TARGET_HOST, TARGET_PORT))
29
30         # DNS over TCP: prepend 2-byte length (big-endian)
31         tcp.sendall(struct.pack("!H", len(data)) + data)
32
33         # Read the 2-byte response length
34         raw_len = tcp.recv(2)
35         if len(raw_len) < 2:
36             tcp.close()
37             return
38         resp_len = struct.unpack("!H", raw_len)[0]
39
40         # Read the full response
41         response = b""
42         while len(response) < resp_len:
43             chunk = tcp.recv(resp_len - len(response))
44             if not chunk:
45                 break
46             response += chunk
47
48         tcp.close()
49
50         # Send the response back over UDP (strip TCP length prefix)
51         if response:
52             sock.sendto(response, client_addr)
53     except Exception:
54         # Silently drop failed queries (malformed, timeout, etc.)
55         pass
56
57
58 def main() -> None:
59     # Create UDP socket
60     sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
61     sock.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
62
63     try:
64         sock.bind((LISTEN_ADDR, LISTEN_PORT))
65     except PermissionError:
66         print(f"dns-proxy: ERROR need root to bind port {LISTEN_PORT}", flush=True)
67         sys.exit(1)
68     except OSError as e:
69         print(f"dns-proxy: ERROR {e}", flush=True)

```

```

70     sys.exit(1)
71
72     print(f"dns-proxy: listening on UDP {LISTEN_ADDR}:{LISTEN_PORT} TCP {TARGET_HOST}:{TARGET_PORT}", flush=
73         True)
74
75     # Use a thread pool to handle concurrent DNS queries efficiently without thread exhaustion
76     executor = ThreadPoolExecutor(max_workers=64)
77
78     while True:
79         try:
80             data, client_addr = sock.recvfrom(4096) # Support EDNS0 (up to 4096 bytes) to prevent truncation
81             executor.submit(handle_query, data, client_addr, sock)
82         except KeyboardInterrupt:
83             break
84         except Exception:
85             pass
86
87     executor.shutdown(wait=False)
88     sock.close()
89
90 if __name__ == "__main__":
91     main()

```

## 34 scripts/fix-docker-bridge-collision.sh

```

1  #!/bin/bash
2  # scripts/fix-docker-bridge-collision.sh nullexit fix for devref 9.13
3  #
4  # Symptom: host default route points at Docker's bridge gateway (172.17.0.1)
5  # instead of the real Wi-Fi gateway. ALL host internet egress fails. Remote
6  # tailnet clients route through the gateway correctly (their tunnel is
7  # independent of the host's broken route table), so they look healthy while
8  # the host itself cannot reach anything.
9  #
10 # Root cause: Docker's default bridge claims 172.17.0.0/16 by default. When the
11 # host's physical Wi-Fi also lands in the same /16 extremely common on
12 # university networks which hand out 172.16.0.0/12 the kernel
13 # installs Docker's bridge as the default route, and every packet bound for
14 # the internet is silently absorbed by docker0 instead of egressing.
15 #
16 # This script:
17 # 1. Detects the actual collision (Docker's 172.17.0.0/16 vs the host's
18 # Wi-Fi subnet, queried live via the ACTIVE interface not hardcoded en0)
19 # 2. Picks a non-colliding Docker bip using /8-family routing (works for
20 # 172.16.0.0/12 campus networks; naive prefix-match would pick 172.26/24
21 # which still lives INSIDE a /12 of 172.16.0.0/12)
22 # 3. Backs up ~/.colima/default/colima.yaml with an ISO-8601 timestamp
23 # 4. Edits the file in place, idempotent (replace existing docker.bip /
24 # append new docker.bip without disturbing other keys; skip YAML comments)
25 # 5. Restarts Colima so the VM rebuilds with the new bridge subnet
26 # 6. Rebinds Wi-Fi DHCP so the host re-learns its real Wi-Fi gateway
27 # 7. Verifies the new default route is no longer the Docker bridge
28 # 8. Re-runs toggle.sh to bring the gateway back up cleanly
29 # 9. Re-runs scripts/diagnose-host-leak.sh and prints the new verdict
30 #
31 # Usage:
32 # bash scripts/fix-docker-bridge-collision.sh # apply the fix
33 # bash scripts/fix-docker-bridge-collision.sh --dry # show what would change, then exit
34 # bash scripts/fix-docker-bridge-collision.sh --skip-toggle # fix but don't restart gateway
35 # bash scripts/fix-docker-bridge-collision.sh --help # usage
36
37 set -uo pipefail
38 # NOTE: errexit (`set -e`) is intentionally NOT enabled so intermediate
39 # `command || warn` patterns can continue past non-fatal failures.
40 # Every critical failure (colima restart, sed, cp backup, toggle.sh)
41 # uses an explicit `|| die` so a half-applied state never escapes.
42
43 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
44 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
45 LOG_FILE="$PROJECT_ROOT/output.log"
46 source "$SCRIPT_DIR/common.sh"
47 TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
48 COLIMA_YAML="$HOME/.colima/default/colima.yaml"
49 PLATFORM="$(uname -s)"
50

```

```

51 # Argument parsing
52 DRY_RUN=false
53 SKIP_TOGGLE=false
54 case "${1:-}" in
55     --dry|-n)      DRY_RUN=true ;;
56     --skip-toggle) SKIP_TOGGLE=true ;;
57     --help|-h)
58         sed -n '28,37p' "$0"
59         exit 0
60     ;;
61     *)
62         : # default: apply
63         ;;
64 *)
65     printf 'Unknown argument: %s\n' "$1" >&2
66     exit 2
67     ;;
68 esac
69
70
71
72 # Helper: print "[dry-run] $cmd" or actually run $cmd. Wraps the redirect
73 # INSIDE the function so outer `>> $FILE` statements never accidentally
74 # append dry-run noise to the real file. (Earlier draft didn't and the
75 # dry-run polluted $COLIMA_YAML.)
76 run() {
77     if [ "$DRY_RUN" = "true" ]; then
78         printf '    [dry-run] %s\n' "$*"
79         return 0
80     fi
81     printf '    $ %s\n' "$*"
82     "$@"
83 }
84
85 # Platform-aware sed in-place. macOS BSD sed needs an empty argument;
86 # GNU sed (Linux) does not. Earlier draft used `-i ''` everywhere and
87 # would have failed on Linux silently.
88 sed_inplace() {
89     local expr="$1"
90     local file="$2"
91     case "$PLATFORM" in
92         Darwin) run sed -i '' "$expr" "$file" ;;
93         *)      run sed -i "$expr" "$file" ;;
94     esac
95 }
96
97 # 0. Resolve the ACTIVE physical interface once (used everywhere)
98 # Replaces the earlier `en0` hardcoding. Same algorithm recover.sh uses:
99 # `route get default` if utun/tun, walk en0..en5 first with `status: active`
100 # + `inet` fall back to en0. This means the script also works on hosts
101 # whose primary NIC is en1/en2 (Thunderbolt Ethernet, USB-LAN, etc.).
102 resolve_active_iface() {
103     local iface
104     iface=$(route get default 2>> "$LOG_FILE" | awk '/interface:/{print $2; exit}')
105     if [ -z "$iface" ] || [ "${iface#utun}" != "$iface" ] || [ "${iface#tun}" != "$iface" ]; then
106         for i in en0 en1 en2 en3 en4 en5; do
107             if ifconfig "$i" 2>> "$LOG_FILE" | grep -q "status: active" \
108                 && ifconfig "$i" 2>> "$LOG_FILE" | grep -q "inet "; then
109                 iface="$i"; break
110             fi
111         done
112     fi
113     [ -z "$iface" ] && iface="en0"
114     printf '%s\n' "$iface"
115 }
116 ACTIVE_IFACE=$(resolve_active_iface)
117
118 # 1. Detect collision
119 step "1/9 Detecting Docker-bridge vs Wi-Fi subnet collision"
120
121 # Host subnet on the ACTIVE interface (not hardcoded en0).
122 HOST_SUBNET=$(ifconfig "$ACTIVE_IFACE" 2>> "$LOG_FILE" \
123     | awk '/inet /{print $2; exit}')
124 HOST_GATEWAY=$(route get default 2>> "$LOG_FILE" \
125     | awk '/gateway:/{print $2; exit}')
126 DEFAULT_DEV=$(route get default 2>> "$LOG_FILE" \
127     | awk '/interface:/{print $2; exit}')
128
129 printf '    active iface: %s\n' "$ACTIVE_IFACE"
130 printf '    iface inet: %s\n' "${HOST_SUBNET:-none}"
131 printf '    default via: %s\n' "${HOST_GATEWAY:-none}"

```

```

132 printf ' default dev: %s\n' "${DEFAULT_DEV:-none}"
133
134 # Collision fires when the host's default-route gateway is 172.17.0.1 (Docker's
135 # bridge) OR when the host subnet is in 172.17.0.0/16. The full-bridge-collision
136 # case (gateway is literally the docker bridge IP) is the smoking gun.
137 COLLISION=false
138 if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#172.17.}" != "$HOST_GATEWAY" ]; then
139     COLLISION=true
140     fail "default gateway $HOST_GATEWAY collides with Docker's default bridge (172.17.0.0/16)"
141 elif [ -n "$HOST_SUBNET" ] && [ "${HOST_SUBNET#172.17.}" != "$HOST_SUBNET" ]; then
142     COLLISION=true
143     fail "host subnet $HOST_SUBNET overlaps Docker's 172.17.0.0/16 bridge"
144 else
145     ok "no collision detected between host Wi-Fi and Docker's default bridge"
146 fi
147
148 if [ "$COLLISION" = "false" ]; then
149     echo ""
150     echo "Nothing to fix. If you still see egress issues, run:"
151     echo "    bash scripts/diagnose-host-leak.sh"
152     exit 0
153 fi
154
155 # 2. Pick non-colliding bip via /8-family routing
156 step "2/9 Picking a non-colliding Docker bip for ~/.colima/default/colima.yaml"
157
158 # Earlier draft tried naive prefix-match: "if candidate's 3-octet prefix is
159 # not a literal prefix of host gateway, pick it". That breaks on /12 campus
160 # networks (172.16.0.0/12 covers 172.16-172.31): a candidate like
161 # 172.26.0.1/24 LIVES INSIDE the host's /12 even though the prefix test
162 # passes. Fix: classify the host's address family FIRST, then choose ALL
163 # candidates from a DIFFERENT RFC1918 family. /8 boundaries are coarse
164 # enough that no /24 picked this way can ever collide.
165 case "$HOST_GATEWAY" in
166     172.16.*|172.17.*|172.18.*|172.19.*|172.20.*|172.21.*|172.22.*|172.23.*|\
172.24.*|172.25.*|172.26.*|172.27.*|172.28.*|172.29.*|172.30.*|172.31.*)
168         HOST_FAMILY="172"
169         # Host is in 172.x jump past 172.31 entirely. Pick from 10.x / 192.168.x.
170         CANDIDATES="10.200.0.1/24 10.201.0.1/24 192.168.99.1/24"
171         ;;
172     10.*)
173         HOST_FAMILY="10"
174         # Host is in 10.x pick from 172.x (well outside any commonly-deployed
175         # 10.0.0.0/8 slice) / 192.168.x.
176         CANDIDATES="172.26.0.1/24 172.28.0.1/24 192.168.99.1/24"
177         ;;
178     192.168.*)
179         HOST_FAMILY="192"
180         CANDIDATES="172.26.0.1/24 10.200.0.1/24 10.201.0.1/24"
181         ;;
182     *)
183         HOST_FAMILY="unknown"
184         warn "host gateway $HOST_GATEWAY is not in any RFC1918 range (/8 family)"
185         warn "falling back to abstract candidate set; verify collision-free manually"
186         CANDIDATES="172.26.0.1/24 10.200.0.1/24 192.168.99.1/24"
187         ;;
188 esac
189
190 printf ' host address family: %s\n' "$HOST_FAMILY"
191 printf ' candidate pool: %s\n' "$CANDIDATES"
192
193 # Sanity-check the chosen pool has at least one entry that doesn't literally
194 # share a /24 with the host gateway. Naive check (just /24 prefix) is fine
195 # here because we already picked across families.
196 CHOSEN=""
197 for cand in $CANDIDATES; do
198     PREFIX3=$(printf '%s' "$cand" | cut -d. -f1-3)
199     if [ -n "$HOST_GATEWAY" ] && [ "${HOST_GATEWAY#$PREFIX3.}" != "$HOST_GATEWAY" ]; then
200         warn "skip $cand (literal /24 collision with host gateway $HOST_GATEWAY)"
201         continue
202     fi
203     CHOSEN="$cand"
204     ok "chose bip=$CHOSEN (cross-family, no /24 literal collision)"
205     break
206 done
207
208 if [ -z "$CHOSEN" ]; then
209     CHOSEN="172.26.0.1/24"
210     warn "could not find a clean candidate; defaulting to $CHOSEN"
211     warn "verify with 'route get default' after restart"
212 fi

```

```

213
214 # 3. Backup colima.yaml
215 step "3/9 Backing up $COLIMA_YAML"
216 mkdir -p "$(dirname "$COLIMA_YAML")"
217 BACKUP="$HOME/.colima/default/colima.yaml.bak.$TIMESTAMP"
218 if [ -f "$COLIMA_YAML" ]; then
219     # `cp` is safe in dry-run because run() suppresses the inner command
220     # entirely no chance of accidentally writing a real backup.
221     run cp -p "$COLIMA_YAML" "$BACKUP" || die "backup failed"
222     if [ "$DRY_RUN" = "false" ]; then ok "backup: $BACKUP"; fi
223 else
224     warn "no existing $COLIMA_YAML (will create fresh)"
225 fi
226
227 # 4. Edit colima.yaml in place (idempotent)
228 step "4/9 Setting docker.bip=$CHOSEN in colima.yaml (in place)"
229
230 # Resolve the existing docker.bip (if any) without picking up commented-out
231 # YAML keys. awk-based extraction; the negated-comment check (`! /^[[:space:]]*#`)
232 # ensures `# bip: 1.2.3.4/24` is treated as a comment, not the active value.
233 EXISTING_BIP=""
234 if [ -f "$COLIMA_YAML" ]; then
235     EXISTING_BIP=$(awk '
236         /^[[:space:]]*#/{next}
237         /^[[:space:]]*bip:[[:space:]]*/{print $2; exit}
238     ' "$COLIMA_YAML" 2>> "$LOG_FILE" || true)
239 fi
240
241 if [ -n "$EXISTING_BIP" ] && [ "$EXISTING_BIP" = "$CHOSEN" ]; then
242     ok "colima.yaml already has docker.bip=$CHOSEN no edit needed"
243 elif [ -n "$EXISTING_BIP" ]; then
244     printf ' (replacing existing docker.bip=%s with %s)\n' "$EXISTING_BIP" "$CHOSEN"
245     # Use sed_inplace so Linux / macOS get the right syntax. Skip YAML comment
246     # lines so existing `# bip:` comments aren't rewired.
247     sed_inplace "s|~\([[:space:]]*bip:[[:space:]]*\)\).*\|~\1$CHOSEN|" "$COLIMA_YAML" \
248         || die "sed replacement failed (colima.yaml malformed?)"
249     ok "colima.yaml updated"
250 else
251     # No existing `bip:` key. Branch:
252     # - file has `docker:` block append ` bip:` under it
253     # - file has no `docker:` block append new docker: section
254     # - file does not exist create fresh
255     # Each branch is dry-run-aware so no branch can accidentally leak noise
256     # into $COLIMA_YAML during a preview.
257     HAS_DOCKER_BLOCK=false
258     if [ -f "$COLIMA_YAML" ] && grep -q '^docker:' "$COLIMA_YAML"; then
259         HAS_DOCKER_BLOCK=true
260     fi
261
262     if [ "$HAS_DOCKER_BLOCK" = "true" ]; then
263         if [ "$DRY_RUN" = "true" ]; then
264             printf '[dry-run] would append to %s:\n bip: %s\n' "$COLIMA_YAML" "$CHOSEN"
265         else
266             printf '\n bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
267                 || die "append failed"
268             ok "appended bip under existing docker: block"
269         fi
270     elif [ -f "$COLIMA_YAML" ]; then
271         if [ "$DRY_RUN" = "true" ]; then
272             printf '[dry-run] would append to %s:\ndocker:\n bip: %s\n' "$COLIMA_YAML" "$CHOSEN"
273         else
274             printf '\ndocker:\n bip: %s\n' "$CHOSEN" >> "$COLIMA_YAML" \
275                 || die "append failed"
276             ok "appended new docker: block at end of file"
277         fi
278     else
279         if [ "$DRY_RUN" = "true" ]; then
280             printf '[dry-run] would create %s with:\ndocker:\n bip: %s\n' "$COLIMA_YAML" "$CHOSEN"
281         else
282             printf 'docker:\n bip: %s\n' "$CHOSEN" > "$COLIMA_YAML" \
283                 || die "create failed"
284             ok "created fresh $COLIMA_YAML with docker.bip=$CHOSEN"
285         fi
286     fi
287 fi
288
289 # Always show the final colima.yaml docker section for visibility
290 printf '\n%b colima.yaml `docker:` section now reads:%b\n' "$BOLD" "$NC"
291 awk '
292     /^docker:/{p=1}
293     p && /^[^ ]/ && !/^docker:/{exit}

```



```

294 p
295 ' "$COLIMA_YAML" 2>/dev/null | sed 's/~ / /' || true
296
297 if [ "$DRY_RUN" = "true" ]; then
298     printf '\n %b(dry-run: exiting before colima restart / dhcp rebind / toggle.sh)%b\n' "$YELLOW" "$NC"
299     exit 0
300 fi
301
302 # 5. Restart Colima
303 step "5/9 Restarting Colima VM with new bip"
304 if ! command -v colima >> "$LOG_FILE" 2>&1; then
305     die "colima not on PATH install with 'brew install colima' first"
306 fi
307 # No timeout wrapper (would mask actual VM startup time). If colima hangs,
308 # the user kills with Ctrl+C and can manually `colima restart` from a fresh
309 # terminal the rest of this script is idempotent so re-running works.
310 colima restart >> "$LOG_FILE" 2>&1 || die "colima restart failed; check $LOG_FILE"
311 ok "colima restarted"
312
313 # 6. Rebind DHCP on the ACTIVE interface (not hardcoded en0)
314 step "6/9 Rebinding DHCP on $ACTIVE_IFACE so host re-learns its real gateway"
315
316 # Map ifname macOS network service name (Wi-Fi, USB 10/100 LAN, etc.).
317 # Same logic as the diagnostic + toggle.sh.
318 ACTIVE_SVC=""
319 if [ "$PLATFORM" = "Darwin" ]; then
320     ACTIVE_SVC=$(get_active_service)
321 fi
322
323 case "$PLATFORM" in
324     Darwin)
325         sudo -n networksetup -setdhcp "$ACTIVE_SVC" renew >> "$LOG_FILE" 2>&1 \
326             && ok "DHCP rebind issued on $ACTIVE_SVC" \
327             || warn "setdhcp renew failed; try 'sudo ipconfig set $ACTIVE_IFACE DHCP' manually"
328         ;;
329     Linux)
330         sudo -n dhclient -r "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 || true
331         sudo -n dhclient "$ACTIVE_IFACE" >> "$LOG_FILE" 2>&1 \
332             && ok "dhclient rebind on $ACTIVE_IFACE" \
333             || warn "dhclient rebind failed"
334         ;;
335     *)
336         warn "unsupported platform $PLATFORM; skip DHCP rebind"
337         ;;
338 esac
339
340 # 7. Verify new default route
341 step "7/9 Verifying the new default route is NOT the Docker bridge"
342 sleep 2 # give DHCP + kernel a beat
343 NEW_DEFAULT=$(netstat -rn 2>> "$LOG_FILE" | awk '/^default/ {print $0; exit}')
344 printf ' %s\n' "${NEW_DEFAULT:-no default route detected}"
345 NEW_GATEWAY=$(printf '%s' "$NEW_DEFAULT" | awk '{print $2; exit}')
346
347 if [ -n "$NEW_GATEWAY" ] && [ "${NEW_GATEWAY#172.17.}" != "$NEW_GATEWAY" ]; then
348     fail "default route STILL points at 172.17.x.x DHCP rebind may need manual touch"
349     warn "try: sudo ipconfig set $ACTIVE_IFACE DHCP (or reconnect Wi-Fi in the menu bar)"
350 elif [ -n "$NEW_GATEWAY" ]; then
351     ok "default route is now via $NEW_GATEWAY (no longer Docker's bridge)"
352 else
353     warn "could not determine new gateway; verify with 'route get default'"
354 fi
355
356 # 8. Re-run toggle.sh to bring the gateway back up
357 if [ "$SKIP_TOGGLE" = "true" ]; then
358     step "8/9 Skipping toggle.sh (--skip-toggle)"
359     warn "gateway is currently down run ./toggle.sh manually when ready"
360 else
361     step "8/9 Re-running toggle.sh to bring the gateway back up"
362     if bash "$SCRIPT_DIR/diagnose-host-leak.sh" 2>&1 | tee -a "$LOG_FILE"; then
363         ok "toggle.sh completed"
364     else
365         warn "toggle.sh returned non-zero see $LOG_FILE for details"
366     fi
367 fi
368
369 # 9. Re-run the diagnostic + show verdict
370 step "9/9 Re-running host-leak diagnostic to confirm fix"
371 if bash "$SCRIPT_DIR/diagnose-host-leak.sh" 2>&1 | tee -a "$LOG_FILE"; then
372     LATEST_REPORT=$(ls -t "$PROJECT_ROOT"/host-leak-diagnostic-*.txt 2>/dev/null | head -1)
373     if [ -n "$LATEST_REPORT" ]; then
374         VERDICT=$(awk '/^ Scenario: / {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)

```

```

375 WARP=$(awk -F': *' '/^ WARP egress:/ {print $2; exit}' "$LATEST_REPORT" 2>/dev/null)
376 printf '\n Most recent verdict: Scenario=%s WARP=%s\n' \
377        "${VERDICT:-?}" "${WARP:-?}"
378 if [ "$VERDICT" = "OK" ]; then
379     ok "fix confirmed by diagnostic host traffic is going through WARP"
380 fi
381 fi
382 else
383     warn "diagnostic exited non-zero; verify manually"
384 fi
385
386 printf '\n\b\b\n' "$BOLD" "$NC"
387 printf '\bDone.\b Full transcript at %s\n' "$BOLD" "$NC" "$LOG_FILE"
388 printf '\b\b\n' "$BOLD" "$NC"
389 exit 0

```

## 35 scripts/fix-ssh-delay.sh

```

1  #!/bin/bash
2  # scripts/fix-ssh-delay.sh Fix ~20s SSH connection delay caused by UseDNS
3  #
4  # macOS OpenSSH defaults UseDNS to "yes". When Tailscale SSH connects from
5  # a peer (phone/laptop), sshd performs a reverse-DNS (PTR) lookup on the
6  # client's 100.x.x.x Tailscale IP. Since DNS is hijacked to the nullexit
7  # gateway AdGuard Cloudflare WARP, and 100.64.0.0/10 is CGNAT private
8  # space, the PTR query times out (~15-20s) before SSH proceeds.
9  #
10 # This script uncomments "UseDNS no" in /etc/ssh/sshd_config and reloads
11 # sshd. Existing SSH sessions are NOT dropped.
12 #
13 # Usage:
14 #   sudo bash scripts/fix-ssh-delay.sh
15 #
16 set -euo pipefail
17
18 SSHD_CONFIG="/etc/ssh/sshd_config"
19
20 # Check we're root
21 if [ "$(id -u)" -ne 0 ]; then
22     echo "ERROR: This script must be run with sudo (it edits $SSHD_CONFIG)"
23     echo "Usage: sudo bash scripts/fix-ssh-delay.sh"
24     exit 1
25 fi
26
27 # Apply config (idempotent)
28 if grep -qE '^UseDNS no' "$SSHD_CONFIG" 2>/dev/null; then
29     echo "UseDNS is already set to 'no' in $SSHD_CONFIG"
30 elif grep -qE '^#UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
31     echo "Uncommenting 'UseDNS no' in $SSHD_CONFIG"
32     sed -i 's/^#UseDNS no/UseDNS no/' "$SSHD_CONFIG"
33 elif grep -qEi '^UseDNS' "$SSHD_CONFIG" 2>/dev/null; then
34     echo "Changing existing UseDNS line to 'UseDNS no'"
35     sed -i 's/^UseDNS.*/UseDNS no/' "$SSHD_CONFIG"
36 else
37     echo "Appending 'UseDNS no' to $SSHD_CONFIG"
38     echo "" >> "$SSHD_CONFIG"
39     echo "UseDNS no" >> "$SSHD_CONFIG"
40 fi
41
42 # Verify
43 if ! grep -qE '^UseDNS no' "$SSHD_CONFIG"; then
44     echo "Failed to apply UseDNS no check $SSHD_CONFIG manually"
45     exit 1
46 fi
47
48 # Reload sshd (does NOT drop existing connections)
49 echo "Reloading sshd..."
50
51 SSHD_PID=$(pgrep -f /usr/sbin/sshd 2>/dev/null | head -1 || true)
52 if [ -n "$SSHD_PID" ]; then
53     kill -HUP "$SSHD_PID"
54     echo "Sent SIGHUP to sshd (PID $SSHD_PID) config reloaded"
55 else
56     echo "No running sshd found. macOS starts sshd on-demand per connection."
57     echo "The new 'UseDNS no' config will take effect on the next SSH connection."
58 fi

```

```

59
60 echo ""
61 echo " Done. SSH connections should now connect instantly."
62 echo " Existing sessions were not interrupted."

```

## 36 scripts/generate\_\_tex.py

```

1  #!/usr/bin/env python3
2  import os
3
4  PROJECT_ROOT = os.path.abspath(os.path.join(os.path.dirname(__file__), '..'))
5
6  IGNORE_DIRS = {'.git', '.github', 'adguard', '__pycache__', 'Recover Gateway.app', 'Toggle Gateway.app', 'Sweep
    Gateway.app', 'nullexit-knowledge-graph'}
7  IGNORE_FILES = {'.env', 'nullexit_unified.tex', 'nullexit_unified.pdf', 'output.log', 'blocked.log', '.DS_Store
    ', '.lan_p2p_detected', '.signatures', '.gateway_ip', 'TUNNEL_FAILED_CLOSED.marker', '.host_ips', 'refactor.
    md', 'sweep.md', '.build_hash', '.rules_hash', 'knowledge-graph.html'}
8  IGNORE_EXTS = {'.pdf', '.tex', '.png', '.jpg', '.jpeg', '.zip', '.tar', '.gz', '.log', '.aux', '.fls', '.
    fdb_latexmk', '.out', '.toc', '.xdv', '.synctex(busy)'}
9
10 # Exclude from code block inclusion, but keep in tree
11 SKIP_CONTENT_FILES = {'LICENSE'}
12
13 PRIORITY_FILES = [
14     'README.md',
15     'agent.md',
16     'devref.md',
17     'diagrams.md',
18     'toggle.sh',
19     'recover.sh',
20     'setup.sh',
21     'scripts/setup-linux.sh',
22     'docker-compose.yml',
23     'scripts/common.sh',
24     'scripts/watcher.sh',
25     'scripts/crypto.sh'
26 ]
27
28 def get_language(filename):
29     ext = os.path.splitext(filename)[1].lower()
30     if ext == '.sh' or ext == '.applescript':
31         return 'bash'
32     elif ext == '.py':
33         return 'Python'
34     elif ext == '.go':
35         return 'Go'
36     elif ext in {'.yaml', '.yml'}:
37         return 'bash'
38     elif ext == '.plist':
39         return 'XML'
40     elif ext == '.md':
41         return ''
42     return ''
43
44 def escape_tex(text):
45     return text.replace('_', '\\_')
46
47 def generate_tree(dir_path, prefix=""):
48     tree_str = ""
49     entries = sorted(os.listdir(dir_path))
50     entries = [e for e in entries if e not in IGNORE_DIRS and e not in IGNORE_FILES and os.path.splitext(e)[1].
        lower() not in IGNORE_EXTS]
51
52     for i, entry in enumerate(entries):
53         path = os.path.join(dir_path, entry)
54         is_last = (i == len(entries) - 1)
55         connector = "-- " if is_last else "|-- "
56         tree_str += f"{prefix}{connector}{entry}\n"
57
58         if os.path.isdir(path):
59             extension = " " if is_last else "| "
60             tree_str += generate_tree(path, prefix + extension)
61
62     return tree_str
63
64 def main():

```



```

144     f.write(tree)
145     f.write(midamble)
146
147     for relpath in all_relpaths:
148         lang = get_language(relpath)
149         lang_attr = f"[language={lang}]" if lang else ""
150
151         f.write(f"\\section{{{escape_tex(relpath)}}}\\n")
152         f.write(f"\\begin{{{lstlisting}}}{lang_attr}\\n")
153         try:
154             with open(os.path.join(PROJECT_ROOT, relpath), 'r', encoding='utf-8', errors='ignore') as src:
155                 content = src.read()
156                 import re
157                 # Strip all non-ASCII characters (emojis, box drawings, em dashes, etc)
158                 # AND ASCII control characters except \t\n (binary payloads, ANSI escapes,
159                 # form feeds all invisible, all fatal to LaTeX compilation)
160                 content = re.sub(r'[\x09\x0A\x20-\x7E]+', '', content)
161                 f.write(content)
162                 if not content.endswith('\n'):
163                     f.write('\n')
164             except Exception as e:
165                 f.write(f"Error reading file: {e}\\n")
166             f.write("\\end{lstlisting}\\n\\n")
167
168         f.write("\\end{document}\\n")
169
170     print(f"Generated {tex_path}")
171
172 if __name__ == '__main__':
173     main()

```

## 37 scripts/logger/go.mod

```

1 module nullexit/logger
2
3 go 1.22

```

## 38 scripts/logger/main.go

```

1 package main
2
3 import (
4     "bufio"
5     "encoding/json"
6     "fmt"
7     "io"
8     "log"
9     "os"
10    "regexp"
11    "strings"
12    "time"
13 )
14
15 const (
16     queryLogPath = "/adguard_work/data/querylog.json"
17     procKmsgPath = "/proc/kmsg"
18     outputLogPath = "/app/blocked.log"
19 )
20
21 func logMessage(msg string) {
22     timestamp := time.Now().Format("2006-01-02 15:04:05")
23     formattedMsg := fmt.Sprintf("%s %s\\n", timestamp, msg)
24     fmt.Print(formattedMsg)
25
26     f, err := os.OpenFile(outputLogPath, os.O_APPEND|os.O_CREATE|os.O_WRONLY, 0644)
27     if err != nil {
28         log.Printf("Error writing to log file: %v\\n", err)
29         return
30     }
31     defer f.Close()
32     f.WriteString(formattedMsg)
33 }

```

```

34
35 func tailFile(filepath string, out chan<- string) {
36     for {
37         if _, err := os.Stat(filepath); os.IsNotExist(err) {
38             time.Sleep(1 * time.Second)
39             continue
40         }
41         break
42     }
43
44     f, err := os.Open(filepath)
45     if err != nil {
46         log.Printf("Error opening file %s: %v\n", filepath, err)
47         return
48     }
49     defer func() {
50         if f != nil {
51             f.Close()
52         }
53     }()
54
55     // Seek to end
56     f.Seek(0, io.SeekEnd)
57     reader := bufio.NewReader(f)
58
59     for {
60         line, err := reader.ReadString('\n')
61         if err != nil {
62             if err == io.EOF {
63                 stat, err := os.Stat(filepath)
64                 if err == nil {
65                     currentOffset, _ := f.Seek(0, io.SeekCurrent)
66                     if stat.Size() < currentOffset {
67                         // File truncated or rotated
68                         f.Close()
69                         for {
70                             if _, err := os.Stat(filepath); os.IsNotExist(err) {
71                                 time.Sleep(1 * time.Second)
72                                 continue
73                             }
74                             break
75                         }
76                         f, _ = os.Open(filepath)
77                         reader = bufio.NewReader(f)
78                         continue
79                     }
80                 }
81                 time.Sleep(500 * time.Millisecond)
82                 continue
83             }
84             log.Printf("Error reading file %s: %v\n", filepath, err)
85             time.Sleep(1 * time.Second)
86             continue
87         }
88         out <- line
89     }
90 }
91
92 type adguardLogEntry struct {
93     IP      string `json:"IP"`
94     QH      string `json:"QH"`
95     QT      string `json:"QT"`
96     Result struct {
97         IsFiltered bool `json:"IsFiltered"`
98         Reason      int  `json:"Reason"`
99         Rules       []struct {
100             Text string `json:"Text"`
101         } `json:"Rules"`
102     } `json:"Result"`
103 }
104
105 func monitorDNS() {
106     logMessage("[System] Starting DNS block logger...")
107     lines := make(chan string)
108     go tailFile(queryLogPath, lines)
109
110     reasonMap := map[int]string{
111         2: "CustomRule",
112         3: "BlockList",
113         4: "SafeBrowsing",
114         5: "ParentalControl",

```

```

115     6: "SafeSearch",
116     7: "BlockedService",
117 }
118
119 for line := range lines {
120     var data adguardLogEntry
121     if err := json.Unmarshal([]byte(line), &data); err != nil {
122         continue
123     }
124
125     res := data.Result
126     if res.IsFiltered || (res.Reason != 0 && res.Reason != 1) {
127         qh := data.QH
128         if qh == "" {
129             qh = "unknown"
130         }
131         qt := data.QT
132         if qt == "" {
133             qt = "unknown"
134         }
135         clientIP := data.IP
136         if clientIP == "" {
137             clientIP = "unknown"
138         }
139
140         ruleText := "unknown rule"
141         if len(res.Rules) > 0 && res.Rules[0].Text != "" {
142             ruleText = res.Rules[0].Text
143         }
144
145         reasonStr, ok := reasonMap[res.Reason]
146         if !ok {
147             reasonStr = fmt.Sprintf("Reason-%d", res.Reason)
148         }
149
150         logMessage(fmt.Sprintf("[DNS] Blocked %s (Type: %s) for client %s | Reason: %s | Rule: %s", qh, qt,
151             clientIP, reasonStr, ruleText))
152     }
153 }
154
155 func monitorIPs() {
156     logMessage("[System] Starting IP block logger...")
157
158     prefixRe := regexp.MustCompile(`(IP_BLOCK_[A-Z_]+):`)
159     srcRe := regexp.MustCompile(`SRC=([0-9a-fA-F\.:]+)`)
160     dstRe := regexp.MustCompile(`DST=([0-9a-fA-F\.:]+)`)
161     protoRe := regexp.MustCompile(`PROTO=([A-Z0-9]+)`)
162     sptRe := regexp.MustCompile(`SPT=(\d+)`)
163     dptRe := regexp.MustCompile(`DPT=(\d+)`)
164     inRe := regexp.MustCompile(`IN=([a-zA-Z0-9\.\_-]+)`)
165     outRe := regexp.MustCompile(`OUT=([a-zA-Z0-9\.\_-]+)`)
166
167     f, err := os.Open(procKmsgPath)
168     if err != nil {
169         if os.IsPermission(err) {
170             logMessage("[System] ERROR: Insufficient permissions to read /proc/kmsg. Enable CAP_SYSLOG.")
171         } else {
172             logMessage(fmt.Sprintf("[System] ERROR in IP logger: %v", err))
173         }
174         return
175     }
176     defer f.Close()
177
178     reader := bufio.NewReader(f)
179     for {
180         line, err := reader.ReadString('\n')
181         if err != nil {
182             if err == io.EOF {
183                 time.Sleep(100 * time.Millisecond)
184                 continue
185             }
186             logMessage(fmt.Sprintf("[System] ERROR reading kmsg: %v", err))
187             time.Sleep(1 * time.Second)
188             continue
189         }
190
191         if strings.Contains(line, "IP_BLOCK") {
192             match := prefixRe.FindStringSubmatch(line)
193             if len(match) > 1 {
194                 prefix := match[1]

```



```

195     src := "unknown"
196     if m := srcRe.FindStringSubmatch(line); len(m) > 1 {
197         src = m[1]
198     }
199
200     dst := "unknown"
201     if m := dstRe.FindStringSubmatch(line); len(m) > 1 {
202         dst = m[1]
203     }
204
205     proto := "unknown"
206     if m := protoRe.FindStringSubmatch(line); len(m) > 1 {
207         proto = m[1]
208     }
209
210     spt := ""
211     if m := sptRe.FindStringSubmatch(line); len(m) > 1 {
212         spt = m[1]
213     }
214
215     dpt := ""
216     if m := dptRe.FindStringSubmatch(line); len(m) > 1 {
217         dpt = m[1]
218     }
219
220     inIf := ""
221     if m := inRe.FindStringSubmatch(line); len(m) > 1 {
222         inIf = m[1]
223     }
224
225     outIf := ""
226     if m := outRe.FindStringSubmatch(line); len(m) > 1 {
227         outIf = m[1]
228     }
229
230     direction := "inbound"
231     if strings.HasSuffix(prefix, "DST") {
232         direction = "outbound"
233     }
234
235     listType := "MALICIOUS"
236     if !strings.Contains(prefix, "MALICIOUS") {
237         parts := strings.Split(prefix, "_")
238         if len(parts) >= 3 {
239             listType = "GEO_" + parts[2]
240         }
241     }
242
243     portStr := ""
244     if dpt != "" {
245         portStr = ":" + dpt
246     }
247
248     srcPortStr := ""
249     if spt != "" {
250         srcPortStr = ":" + spt
251     }
252
253     ifStr := ""
254     if inIf != "" && outIf != "" {
255         ifStr = fmt.Sprintf("IF: %s->%s", inIf, outIf)
256     } else if inIf != "" || outIf != "" {
257         ifStr = fmt.Sprintf("IF: %s%s", inIf, outIf)
258     }
259
260     logMessage(fmt.Sprintf("[IP] Blocked %s to %s%s (Proto: %s) from %s%s (%s) | List: %s", direction, dst,
261         portStr, proto, src, srcPortStr, ifStr, listType))
262 }
263 }
264 }
265 }
266
267 func main() {
268     logMessage("[System] Nullexit Block Logger starting...")
269
270     go monitorDNS()
271     monitorIPs()
272 }

```

## 39 scripts/pf.conf

```
1 # scripts/pf.conf - Packet Filter configuration for nullexit killswitch
2 # Dynamically parameterized by pfctl macro override
3
4 # ext_if must be passed via command line, e.g. -D ext_if=en0
5 # If not passed, default to en0
6 ext_if = "en0"
7 ts_if = "utun+"
8
9 # Tables populated dynamically via pfctl
10 table <vpn_endpoints> persist
11 table <derp_relays> persist
12
13 # Core Rules
14 set skip on lo0 # Ensure local loopback is never blocked
15 scrub out all max-mss 1160 # TCP MSS Clamping: Prevents fragmentation on host TCP flows (like Tailscale
16   control plane)
17 pass quick on $ts_if all # Allow all Tailscale traffic (essential to prevent SSH lockout)
18
19 # Default deny outbound WAN traffic
20 block out on $ext_if all
21
22 # Block inbound Screen Sharing (VNC) and SSH from local Wi-Fi, forcing them to use Tailscale
23 block in on $ext_if proto tcp from any to any port { 22, 5900 }
24
25 # Exceptions for tunnel handshakes and coordination
26 pass out quick on $ext_if proto udp to <vpn_endpoints> port 2408
27 pass out quick on $ext_if proto udp to <derp_relays>
28 pass out quick on $ext_if proto tcp to 192.200.0.0/24 port 443
29
30 # Allow local LAN traffic (AirDrop, local Tailscale P2P, etc)
31 pass out quick on $ext_if proto tcp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
32 pass out quick on $ext_if proto udp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
33 pass out quick on $ext_if proto icmp to { 192.168.0.0/16, 10.0.0.0/8, 172.16.0.0/12 }
34
35 # Allow DHCP outbound (client port 68 to server port 67) so we can obtain/renew IP address
36 pass out quick on $ext_if proto udp from any port 68 to any port 67
```

## 40 scripts/reinstall-host-tailscale.sh

```
1 #!/usr/bin/env bash
2 #
3 # scripts/reinstall-host-tailscale.sh
4 #
5 # Complete uninstall + reinstall of the host-side Homebrew Tailscale.
6 #
7 # Use this when the brew LaunchAgent is wedged, Login Items entries are
8 # duplicated, "Tailscale is stopped" persists despite
9 # `brew services restart tailscale`, or you simply want a clean baseline
10 # before re-engaging the gateway as exit node.
11 #
12 # Idempotent. Safe to re-run. Use --dry to preview every step with zero
13 # changes. Use --yes to skip the per-step confirmation prompts.
14 #
15 # This script ONLY touches the HOST-side Tailscale (the one installed by
16 # Homebrew `brew install tailscale` and managed by launchd as
17 # homebrew.mxcl.tailscale). The nullexit gateway container's tailscaled
18 # (which lives inside the Colima VM at 100.100.21.8 and offers the exit
19 # node to this host) is unaffected.
20
21 set -u
22
23 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
24
25 DRY=0
26 YES=0
27
28 for arg in "$@"; do
29   case "$arg" in
30     --dry) DRY=1 ;;
31     --yes) YES=1 ;;
32     --help|-h) # NOTE: heredoc is UNQUOTED so $SCRIPT_DIR expands in the body. If you
33               # add new $-style placeholders here, they'll expand the same way.
34               cat <<USAGE
```

```

35 Usage: bash "$SCRIPT_DIR/reinstall-host-tailscale.sh" [--dry] [--yes]
36
37 --dry Print every step, make zero changes. Use to audit before
38 committing.
39 --yes Auto-confirm the prompts at every destructive step. Useful in
40 CI / fully-scripted recovery; do not use until you've read
41 the script.
42
43 After this script completes cleanly, you still must (MANUALLY):
44
45 1. Open System Settings Privacy & Security Local Network and
46 enable Tailscale. macOS will prompt automatically the first time
47 the fresh install of `tailscale` is invoked.
48
49 2. Re-engage the tailnet + exit node:
50     sudo tailscale up --ssh=true --accept-dns=false \
51       --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{print $1; exit}')" \
52       --exit-node-allow-lan-access=true
53
54 3. Re-run ./toggle.sh from the project root.
55
56 4. Final verification:
57     curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)='
58     expected: warp=on, ip=<Cloudflare>, colo=YYZ/YUL
59
60 USAGE
61     exit 0
62     ;;
63 *)
64     echo "Unknown flag: $arg" >&2
65     exit 2
66     ;;
67 esac
68 done
69
70 # ----- helpers -----
71
72 confirm() {
73     if [ "$YES" = 1 ] || [ "$DRY" = 1 ]; then
74         return 0
75     fi
76     printf "      %s [y/N] " "$1"
77     read -r ans
78     case "$ans" in
79         y|Y|yes|YES) return 0 ;;
80         *)
81             echo "      aborted by user"
82             exit 1
83             ;;
84     esac
85 }
86
87 header() {
88     printf '\n %s\n' "$1"
89 }
90
91 # with_timeout <seconds> <cmd...>
92 # Run a command in the background and kill it if it exceeds <seconds>.
93 # Returns the command's exit code if it finished in time, 124 (matching
94 # GNU `timeout` convention) if it timed out. macOS doesn't ship GNU
95 # coreutils `timeout`, so this is the bash-3.2-friendly replacement.
96 # stdout + stderr are discarded (we only care about exit code); redirect
97 # or pipe externally if you need the output.
98 with_timeout() {
99     local timeout_s="${1:-30}"
100     shift
101     "$@" >/dev/null 2>&1 &
102     local cmdpid=$!
103     local elapsed=0
104     while kill -0 "$cmdpid" 2>/dev/null; do
105         if [ "$elapsed" -ge "$timeout_s" ]; then
106             kill -9 "$cmdpid" 2>/dev/null
107             wait "$cmdpid" 2>/dev/null
108             return 124
109         fi
110         sleep 1
111         elapsed=$((elapsed+1))
112     done
113     wait "$cmdpid" 2>/dev/null
114     return $?
115 }

```

```

116 # ----- 0. Pre-flight -----
117 header "Pre-flight current host tailscale state"
118 echo " brew: $(command -v brew >/dev/null 2>&1 && brew --version | head -1 || echo 'not on
PATH')"
119 echo " tailscaled running: $(pgrep -lf tailscaled 2>/dev/null | wc -l | tr -d ' ') process(es)"
120 echo " brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{print $2, $3}' | tr
'\n' ' ' | sed 's/ $//')"
121 echo " LaunchAgent plists (under ~/Library/LaunchAgents/):"
122 ls -l ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/~/ /' || echo " (none)"
123
124 # Detect Tailscale's macOS Network Extension (residual from prior App Store
125 # Tailscale.app installs). Lives under /Library/Apple/SystemExtensions/ and
126 # is managed by `systemextensionsctl` NOT launchd. If loaded, brew
127 # tailscaled will NOT engage the Network Extension framework slot because
128 # the SE still has it claimed. Symptom: 'sudo tailscale status' returns
129 # "Tailscale is stopped" even when the daemon is alive and listening.
130 se_found=0
131 if command -v systemextensionsctl >/dev/null 2>&1; then
132 # Grab any line matching tailscale and network-extension
133 se_line=$(systemextensionsctl list 2>/dev/null | grep -iE '(io|com)\.tailscale\..*network-extension' | head -
n 1)
134 if [ -n "$se_line" ]; then
135 se_found=1
136 # Extract the teamID (10-char uppercase/numeric word)
137 se_teamID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '[A-Z0-9]{10}$' | head -n 1)
138 # Extract the bundleID (starts with io.tailscale or com.tailscale)
139 se_bundleID=$(echo "$se_line" | tr -s ' \t' '\n' | grep -E '(io|com)\.tailscale\.' | head -n 1 | sed 's
/(.*)//')
140
141 # If teamID is not found, fallback to '-'
142 if [ -z "$se_teamID" ]; then
143 se_teamID="-"
144 fi
145 fi
146 fi
147
148 if [ "$se_found" = 1 ] && [ -n "$se_bundleID" ]; then
149 echo " Tailscale Network Extension System Extension is loaded (residual from App Store Tailscale.app
install)"
150 echo " Team ID: $se_teamID, Bundle ID: $se_bundleID"
151 echo " brew tailscaled cannot engage the NE slot while this SE has it claimed."
152 if [ "$DRY" = 1 ]; then
153 echo " [dry] (would prompt) sudo systemextensionsctl uninstall $se_teamID $se_bundleID"
154 else
155 confirm "Run 'sudo systemextensionsctl uninstall $se_teamID $se_bundleID'? (frees the NE slot for brew
tailscale; non-fatal if it errors)"
156 sudo systemextensionsctl uninstall "$se_teamID" "$se_bundleID" 2>&1 \
|| echo " ! uninstall returned non-zero (may require reboot, non-fatal)"
157 fi
158 fi
159 else
160 echo " no Tailscale System Extension loaded"
161 fi
162
163 # Pattern used by the Step-1 bootout loops, in both dry + real branches.
164 # Widened from a bare /tailscale/ so the same loops also pick up any
165 # nullexit-labeled LaunchAgents (e.g., com.nullexit.wake-recovery, the
166 # caffeinate + post-wake recovery hook). The wake-recovery LaunchAgent
167 # is re-installed by ./toggle.sh / setup.sh on re-engage, so wiping it
168 # here is non-fatal and is part of restoring a clean baseline.
169 TAILSCALE_LABEL_RE='tailscale'
170 NULLEXIT_LABEL_RE='nullexit'
171
172 # ----- 1. brew services stop -----
173 header "Step 1 Stop the brew-managed service + unload LaunchAgents"
174 # Detect whether the brew-managed service is registered as $USER (gui/$UID
175 # domain) or as root (system domain). NOPASSWD sudoers can leave a stale
176 # root-owned registration after a prior 'sudo brew services ...' round;
177 # 'brew services stop' without sudo refuses to touch a root-owned plist
178 # (Error: Service 'tailscale' is started as `root`).
179 brew_status_line=$(brew services list 2>/dev/null | awk '$1=="tailscale"')
180 brew_status_user=$(echo "$brew_status_line" | awk '{print $3}')
181 if [ -n "$brew_status_line" ]; then
182 if [ "$DRY" = 1 ]; then
183 echo " [dry] brew services stop tailscale (user=$brew_status_user)"
184 else
185 if [ "$brew_status_user" = "root" ]; then
186 confirm "Run 'sudo brew services stop tailscale' (service is registered as root)?"
187 sudo brew services stop tailscale 2>&1 || echo " (sudo brew services stop returned non-zero non-fatal
)"
188 else
189 echo " brew services stop tailscale (user=$brew_status_user, no sudo needed)"

```

```

190     brew services stop tailscale 2>&1 || echo "      (brew reported it wasn't actually running)"
191 fi
192 fi
193 echo "      unload any tailscale-related LaunchAgents (prevents launchd from respawning after kill)"
194 if [ "$DRY" = 1 ]; then
195     echo "      [dry] for each 'tailscale|nullexit'-matched label in gui/$UID domain:"
196     while IFS= read -r label; do
197         [ -z "$label" ] && continue
198         echo "      [dry] launchctl bootout gui/$UID/$label"
199         done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 &&
200         ($3 ~ ts || $3 ~ nx) {print $3}')
201     echo "      [dry] for each 'tailscale|nullexit'-matched label in system domain (root-owned brew plist):"
202     while IFS= read -r label; do
203         [ -z "$label" ] && continue
204         echo "      [dry] (would prompt) sudo launchctl bootout system/$label"
205         done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
206         NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
207         if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
208             echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macos"
209         fi
210         if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
211             echo "      [dry] (would prompt) sudo launchctl bootout system/com.tailscale.ipn.macsys"
212         fi
213     else
214         # Bootout every gui-domain LaunchAgent whose label matches 'tailscale'
215         # OR 'nullexit'. The nullexit arm catches com.nullexit.wake-recovery,
216         # which holds caffeinate/watcher.sh alive across a daemon reset not
217         # desired while the script is establishing a fresh baseline. Awake-
218         # recovery is re-installed by ./toggle.sh / setup.sh on re-engage, so
219         # wiping here is non-fatal. The pattern stays narrower than a bare
220         # /tail/ or /null/ so unrelated Apple services are not unloaded.
221         while IFS= read -r label; do
222             [ -z "$label" ] && continue
223             echo "      launchctl bootout gui/$UID/$label"
224             launchctl bootout "gui/$UID/$label" 2>/dev/null || true
225             done < <(launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" 'NR > 1 &&
226             ($3 ~ ts || $3 ~ nx) {print $3}')
227         # Bootout every system-domain LaunchDaemon/Agent whose label matches
228         # 'tailscale' OR 'nullexit'. This catches the root-owned brew plist left
229         # behind by prior 'sudo brew services ...' invocations that was the
230         # case that left tailscaled PID 47447 57734 respawning between sudo
231         # pkill runs. Defensive nullexit catch in case the user previously
232         # `sudo ln -s`d the wake-recovery plist into /Library/LaunchDaemons.
233         while IFS= read -r label; do
234             [ -z "$label" ] && continue
235             confirm "Run 'sudo launchctl bootout system/$label'?"
236             sudo launchctl bootout "system/$label" 2>/dev/null || true
237             done < <(sudo -n launchctl list 2>/dev/null | awk -v ts="$TAILSCALE_LABEL_RE" -v nx="$NULLEXIT_LABEL_RE" '
238             NR > 1 && ($3 ~ ts || $3 ~ nx) {print $3}')
239         # Legacy Tailscale.app system-domain entries. Each one is checked and
240         # confirmed separately so we don't sudo-prompt or sudo-bootout for nothing.
241         if launchctl print system/com.tailscale.ipn.macos >/dev/null 2>&1; then
242             confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macos'?"
243             sudo launchctl bootout system/com.tailscale.ipn.macos 2>/dev/null || true
244         fi
245         if launchctl print system/com.tailscale.ipn.macsys >/dev/null 2>&1; then
246             confirm "Run 'sudo launchctl bootout system/com.tailscale.ipn.macsys'?"
247             sudo launchctl bootout system/com.tailscale.ipn.macsys 2>/dev/null || true
248         fi
249     fi
250     echo "      waiting up to 10s for tailscaled to exit"
251     for i in 1 2 3 4 5 6 7 8 9 10; do
252         if [ "$DRY" = 1 ]; then
253             echo "      [dry] sleep 1; loop check would terminate here"
254             break
255         fi
256         sleep 1
257         if ! pgrep -lf tailscaled >/dev/null 2>&1; then
258             echo "      tailscaled exited in ${i}s"
259             break
260         fi
261         [ "$i" = 10 ] && echo "      tailscaled still alive after 10s Step 2 will deal with it"
262     done
263 else
264     echo "      brew reports no tailscale service loaded"
265 fi
266
267 # ----- 2. Kill any orphan tailscaled -----
268 header "Step 2 Kill any orphan tailscaled process(es) (may escalate to sudo)"
269 if ! pgrep -lf tailscaled >/dev/null 2>&1; then
270     echo "      no orphans"

```

```

267 else
268     echo "      Stray process(es):"
269     pgrep -lf tailscaled | sed 's/^/      /'
270     echo "      no-sudo: pkill tailscaled"
271     if [ "$DRY" = 1 ]; then
272         echo "      [dry] pkill tailscaled on failure: sudo pkill -9 tailscaled"
273     else
274         if pkill tailscaled 2>/dev/null; then
275             echo "      no-sudo pkill succeeded"
276         else
277             echo "      ! no-sudo pkill insufficient; cannot escalate yet"
278         fi
279         sleep 1
280     fi
281
282     # Recheck; if still alive, offer sudo escalation
283     if { [ "$DRY" = 0 ] && pgrep -lf tailscaled >/dev/null 2>&1; } || [ "$DRY" = 1 ]; then
284         if [ "$DRY" = 1 ]; then
285             echo "      [dry] sudo pkill -9 tailscaled (escalation would fire here)"
286         else
287             confirm "Run 'sudo pkill -9 tailscaled' to escalate?"
288             sudo pkill -9 tailscaled 2>&1
289             sleep 1
290         fi
291     fi
292
293     if [ "$DRY" = 0 ]; then
294         if pgrep -lf tailscaled >/dev/null 2>&1; then
295             echo "      Failed to kill stragglers bailing out"
296             pgrep -lf tailscaled | sed 's/^/      /'
297             exit 3
298         fi
299         echo "      all tailscaled processes gone"
300     fi
301 fi
302
303 # ----- 3. brew uninstall -----
304 header "Step 3 brew uninstall tailscale"
305 if ! brew list tailscale >/dev/null 2>&1; then
306     echo "      not currently installed via brew skip"
307 else
308     # Detect whether the keg directory is owner=$USER or owner=root. If the
309     # keg is root-owned (typically a residual of some prior 'sudo brew ...'
310     # cycle that did perms fixups), 'brew uninstall' without sudo refuses
311     # to remove the files and prints: 'Error: Could not remove tailscale
312     # keg! Do so manually: sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>'.
313     # Detect that and escalate to a sudoed uninstall, with a manual rm -rf
314     # fallback if `sudo brew uninstall` itself errors out (e.g., partial
315     # state where the keg dir is gone but brew still thinks it's installed).
316     # Detect keg state. Treat `(dir absent)` AND `(dir present + empty)` as the
317     # same `<unknown>` bucket so both escalate to sudo. The empty-parent case
318     # is what you get after a manual `sudo rm -rf /opt/homebrew/Cellar/tailscale/<v>`
319     # from inside: brew's parent dir is left present-but-empty and brew still
320     # thinks the package is installed.
321     keg_dir="/opt/homebrew/Cellar/tailscale"
322     if [ -d "$keg_dir" ] && [ -n "$(ls -A "$keg_dir" 2>/dev/null)" ]; then
323         keg_owner="$(stat -f '%Su' "$keg_dir" 2>/dev/null || echo '<unknown>')"
324     else
325         keg_owner="<unknown>"
326     fi
327     if [ "$DRY" = 1 ]; then
328         echo "      [dry] brew uninstall tailscale (keg_owner=$keg_owner)"
329     else
330         # 'root' AND '<unknown>' both require sudo escalation. The '<unknown>'
331         # case is the partial-state path: keg dir was already removed (manually
332         # or by a prior failed uninstall) but brew's internal state still says
333         # installed. Without this branch the script would fall through to a
334         # non-sudo brew uninstall that fails again with the same 'Could not
335         # remove tailscale keg!' error.
336         if [ "$keg_owner" = "root" ] || [ "$keg_owner" = "<unknown>" ]; then
337             confirm "Run 'sudo brew uninstall tailscale' (keg=$keg_owner)?"
338             sudo brew uninstall tailscale 2>&1 || {
339                 echo "      ! sudo brew uninstall failed; offering manual rm -rf fallback"
340                 # Guard the glob so an empty/absent dir doesn't literal-expand the
341                 # '*' and confuse the user with a sudo error about a non-existent path.
342                 if [ -d /opt/homebrew/Cellar/tailscale ] && [ -n "$(ls -A /opt/homebrew/Cellar/tailscale 2>/dev/null)"
343 ]; then
344                     confirm "Run 'sudo rm -rf /opt/homebrew/Cellar/tailscale/*'?"
345                     sudo rm -rf /opt/homebrew/Cellar/tailscale/* 2>&1
346                 else
347                     echo "      keg dir already empty or absent nothing more to rm"

```

```

347     fi
348 }
349 else
350     confirm "Run 'brew uninstall tailscale' (removes LaunchAgent plist + linked symlinks)?"
351     brew uninstall tailscale 2>&1
352 fi
353 fi
354 fi
355
356 # ----- 4. Wipe state directories -----
357 header "Step 4 Wipe stale state"
358 # User-owned
359 declare -a USER_PATHS
360 USER_PATHS=(
361     "$HOME/.local/share/tailscale"
362     "$HOME/.local/run/tailscaled.socket"
363     "$HOME/.local/run/tailscaled.state"
364     "$HOME/.local/state/tailscale"
365     "$HOME/Library/Application Support/Tailscale"
366     "$HOME/Library/Caches/com.tailscale.ipn.macos"
367     "$HOME/Library/Caches/Tailscale"
368     "$HOME/Library/Logs/Tailscale"
369 )
370 echo "    User-owned paths (no sudo):"
371 for p in "${USER_PATHS[@]}; do
372     if [ -e "$p" ]; then
373         if [ "$DRY" = 1 ]; then
374             echo "        [dry] rm -rf '$p'"
375         else
376             rm -rf "$p"
377             echo "        rm -rf '$p'"
378         fi
379     fi
380 done
381
382 # System
383 declare -a SUDO_PATHS
384 SUDO_PATHS=(
385     "/var/lib/tailscale"
386     "/var/cache/tailscale"
387 )
388 echo "    System paths under /var (will prompt for sudo if any are present):"
389 eligible_sudo=0
390 for p in "${SUDO_PATHS[@]}; do
391     if [ -e "$p" ]; then
392         eligible_sudo=1
393         if [ "$DRY" = 1 ]; then
394             echo "        [dry] sudo rm -rf '$p'"
395         else
396             if confirm "Run 'sudo rm -rf $p'?" ; then
397                 sudo rm -rf "$p" && echo "        rm -rf '$p'"
398             fi
399         fi
400     fi
401 done
402 [ "$eligible_sudo" = 0 ] && echo "        (none of /var/lib/tailscale, /var/cache/tailscale are present)"
403
404 # ----- 5. Wipe stale LaunchAgent plists -----
405 header "Step 5 Wipe stale LaunchAgent plist files"
406 declare -a PLIST_FILES
407 PLIST_FILES=(
408     "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
409     "$HOME/Library/LaunchAgents/com.tailscale.ipn.macos.plist"
410     "$HOME/Library/LaunchAgents/homebrew.mxcl.tailscale.plist.bak"
411     "$HOME/Library/LaunchAgents/com.nullexit.wake-recovery.plist"
412 )
413 for f in "${PLIST_FILES[@]}; do
414     if [ -e "$f" ]; then
415         if [ "$DRY" = 1 ]; then
416             echo "        [dry] rm -f '$f'"
417         else
418             rm -f "$f"
419             echo "        rm -f '$f'"
420         fi
421     fi
422 done
423 echo "    Remaining tailscale|nullexit plists (should be empty):"
424 ls -l ~/Library/LaunchAgents/ 2>/dev/null | grep -i 'tail' | sed 's/~/ /' || echo "    (none)"
425
426 # ----- 5.5. Drain Background-Items (Login Items) -----
427 header "Step 5.5 Drain Background-Items (Login Items)"

```



```

428 # macOS 13+ keeps a parallel "Background Items" database alongside
429 # classic LaunchAgents. Even when we wipe ~/Library/LaunchAgents/homebrew.mxcl.tailscale.plist
430 # in Step 5, the .btm file at
431 # ~/Library/Application Support/com.apple.backgroundtaskmanagementagent.backgrounditems.btm
432 # can still hold a Background-Items entry pointing at the (now-dead) plist
433 # path. Next login: macOS sees the entry, tries to bootstrap, finds the
434 # plist missing, and either silently regenerates it via brew's post-install
435 # hook (causing tailscaled respawn at next login) or refuses outright.
436 # `sftool reset-login-items` is Apple's canonical API for wiping the
437 # entire .btm file. It DOES wipe all Login Items the user re-adds what
438 # they want via System Settings General Login Items. Acceptable here
439 # because the goal of this script is "fresh baseline for Tailscale".
440 BTM="$HOME/Library/Application Support/com.apple.backgroundtaskmanagementagent/backgrounditems.btm"
441 if [ -f "$BTM" ]; then
442     if [ "$DRY" = 1 ]; then
443         echo "    [dry] found $BTM; would prompt sftool reset-login-items"
444     else
445         confirm "Run 'sftool reset-login-items' (wipes ALL Login Items; rebuild via System Settings General
446             Login Items)?"
447         if sftool reset-login-items 2>&1; then
448             echo "        Login Items drained"
449         else
450             echo "        ! sftool reset-login-items failed (non-fatal; do manually: System Settings General Login
451                 Items)"
452         fi
453     fi
454 else
455     echo "    no backgrounditems.btm present"
456 fi
457
458 # ----- 6. brew install -----
459 header "Step 6 brew install tailscale (fresh)"
460 if brew list tailscale >/dev/null 2>&1; then
461     echo "    already installed skip"
462 else
463     confirm "Run 'brew install tailscale'?"
464     if [ "$DRY" = 1 ]; then
465         echo "    [dry] brew install tailscale"
466     else
467         brew install tailscale 2>&1
468     fi
469 fi
470
471 # ----- 6.5. Strip quarantine xattr from keg -----
472 header "Step 6.5 Strip quarantine xattr from tailscale install"
473 # brew's freshly poured bottles don't carry com.apple.quarantine (the
474 # bottles are content-addressed and Homebrew-verified), but the upstream
475 # installer (Tailscale.app from App Store) and any manual download do.
476 # Strip the xattr from the entire keg as a defensive safety net so the
477 # freshly-installed tailscaled binary and any helper binaries are not
478 # Gatekeeper-blocked on first launch. Idempotent no-op if no quarantine
479 # xattr is present.
480 quarantine_count="$(xattr -lr /opt/homebrew/Cellar/tailscale 2>/dev/null | grep -c '^[^ ]* com.apple.quarantine
481     ' || echo 0)"
482 if [ "${quarantine_count:-0}" -gt 0 ]; then
483     if [ "$DRY" = 1 ]; then
484         echo "    [dry] found $quarantine_count file(s) with com.apple.quarantine; would prompt xattr -dr"
485     else
486         confirm "Run 'xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale' (strips quarantine xattr from
487             $quarantine_count file(s); non-fatal if it errors)"
488         if xattr -dr com.apple.quarantine /opt/homebrew/Cellar/tailscale 2>&1; then
489             echo "        quarantine xattr stripped"
490         else
491             echo "        ! xattr -dr failed (non-fatal; open /opt/homebrew/Cellar/tailscale in Finder, right-click
492                 tailscaled Open Anyway)"
493         fi
494     fi
495 else
496     echo "    no com.apple.quarantine xattr on tailscale install"
497 fi
498
499 # ----- 7. brew services start (NO sudo!) -----
500 header "Step 7 Start the service as $USER (NO sudo, with multi-tier self-healing)"
501 if ! command -v tailscale >/dev/null 2>&1; then
502     echo "    tailscale not on PATH after install brew install errored; aborting"
503     exit 4
504 fi
505
506 # Aliveness fan-out: ANY one of these is taken as proof of life. None of
507 # them alone is reliable pgrep misses tailscaled if it does setproctitle
508 # or fork-detach; launchctl's reported PID can be a shim that exited

```

```

504 # within 1s; tailscale-cli hangs while macOS Local-Network permission is
505 # unresolved. Combined, they cover each other's blind spots.
506
507 tailscaled_api_up() {
508     # Strategy A: tailscale CLI can talk to the local API. Canonical proof.
509     # Bound at 8s tailscale status itself returns quickly on success OR
510     # failure, but if the post-install state is wedged it can hang.
511     with_timeout 8 /opt/homebrew/bin/tailscale status >/dev/null 2>&1
512 }
513
514 tailscaled_proc_up() {
515     # Strategy B: any process has 'tailscaled' substring in argv.
516     pgrep -lf tailscaled >/dev/null 2>&1
517 }
518
519 tailscaled_lc_up() {
520     # Strategy C: launchctl-managed service has a non-dash, non-zero PID.
521     local lc_pid
522     lc_pid="$(launchctl list 2>/dev/null | awk '$3 == "homebrew.mxcl.tailscale" {print $1; exit}')"
523     if [ -n "$lc_pid" ] && [ "$lc_pid" != "-" ] && [ "$lc_pid" -gt 0 ] 2>/dev/null; then
524         return 0
525     fi
526     return 1
527 }
528
529 tailscaled_is_alive() {
530     tailscaled_api_up && return 0
531     tailscaled_proc_up && return 0
532     tailscaled_lc_up && return 0
533     return 1
534 }
535
536 confirm "Run 'brew services start tailscale'?"
537 if [ "$DRY" = 1 ]; then
538     echo " [dry] brew services start tailscale"
539     echo " [dry] (then poll up to 30s for tailscaled_is_alive)"
540     echo " [dry] (recovery tier 1: launchctl kickstart)"
541     echo " [dry] (recovery tier 2: launchctl bootout + bootstrap of explicit plist)"
542     echo " [dry] (recovery tier 3: brew services restart (stop+start cycle))"
543     echo " [dry] (recovery tier 4: direct /opt/homebrew/bin/tailscaled spawn (bypasses launchd))"
544 else
545     brew services start tailscale 2>&1
546     # Primary poll tailscaled_is_alive covers all 3 detection strategies.
547     echo " polling for tailscaled to come up (up to ~30s)"
548     tailscaled_up=0
549     for i in 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15; do
550         if tailscaled_is_alive; then
551             tailscaled_up=1
552             echo " tailscaled alive after ${i}*2s"
553             break
554         fi
555         sleep 2
556     done
557
558     # Recovery tier 1: kickstart the launchd-managed service. This is the
559     # cheapest and most-aligned-with-launchd fix when brew has loaded the
560     # LaunchAgent but tailscaled didn't fork. Kickstart tells launchd
561     # "re-spawn if not running" without forcing a teardown.
562     if [ "$tailscaled_up" = 0 ]; then
563         echo " ! Recovery tier 1: launchctl kickstart"
564         if launchctl kickstart -kp "gui/$UID/homebrew.mxcl.tailscale" 2>&1; then
565             for i in 1 2 3 4 5 6 7 8 9 10; do
566                 if tailscaled_is_alive; then
567                     tailscaled_up=1
568                     echo " tailscaled alive after kickstart (${i}*2s)"
569                     break
570                 fi
571                 sleep 2
572             done
573         else
574             echo " ! kickstart returned non-zero (LaunchAgent may not be loaded)"
575         fi
576     fi
577
578     # Recovery tier 2: explicit plist bootout + bootstrap. Forces a fresh
579     # launchd registration cycle for the brew-managed service, which fixes
580     # the wider class of "brew says started but daemon never spawned"
581     # wedge states that kickstart alone can't dig out of.
582     if [ "$tailscaled_up" = 0 ]; then
583         echo " ! Recovery tier 2: launchctl bootout + bootstrap of explicit plist"
584         launchctl bootout "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null || true

```

```

585 sleep 1
586 # Brew installs the canonical plist template at $HOMEBREW_PREFIX/Library/.
587 # On Apple Silicon that's /opt/homebrew/Library/LaunchAgents/. On Intel
588 # it's /usr/local/Library/LaunchAgents/. Detect which and pick the
589 # right one so this script works across both architectures.
590 if [ -f /opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
591     plist_path="/opt/homebrew/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
592 elif [ -f /usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist ]; then
593     plist_path="/usr/local/Library/LaunchAgents/homebrew.mxcl.tailscale.plist"
594 else
595     plist_path=""
596 fi
597 if [ -n "$plist_path" ]; then
598     if launchctl bootstrap "gui/$UID" "$plist_path" 2>&1; then
599         for i in 1 2 3 4 5 6 7 8 9 10; do
600             if tailscaled_is_alive; then
601                 tailscaled_up=1
602                 echo "          tailscaled alive after bootstrap (${i}*2s)"
603                 break
604             fi
605             sleep 2
606         done
607     else
608         echo "          ! explicit bootstrap returned non-zero"
609     fi
610 else
611     echo "          ! could not locate brew plist template; skipping bootstrap tier"
612 fi
613 fi
614
615 # Recovery tier 3: brew services restart (forces a stop + start cycle,
616 # which unsticks LaunchAgents with stale 'started' state from prior
617 # half-cycles). This is the canonical brew-doc fix for 'services don't
618 # actually start even though brew says started'.
619 if [ "$tailscaled_up" = 0 ]; then
620     echo "          ! Recovery tier 3: brew services restart (forces stop+start cycle)"
621     if brew services restart tailscale 2>&1; then
622         for i in 1 2 3 4 5 6 7 8 9 10; do
623             if tailscaled_is_alive; then
624                 tailscaled_up=1
625                 echo "          tailscaled alive after brew restart (${i}*2s)"
626                 break
627             fi
628             sleep 2
629         done
630     else
631         echo "          ! brew services restart returned non-zero"
632     fi
633 fi
634
635 # Recovery tier 4: last-resort direct spawn of the daemon binary,
636 # bypassing launchd entirely. If tailscaled crashes here too, we'll
637 # capture its actual stderr in /tmp for the user to read (NOPASSWD
638 # sudoers often have local-network permission denied the captured
639 # stderr makes it obvious what to fix in System Settings).
640 if [ "$tailscaled_up" = 0 ]; then
641     echo "          ! Recovery tier 4: direct spawn (bypasses launchd; tail stderr to /tmp)"
642     /opt/homebrew/bin/tailscaled >/tmp/nullexit-tailscaled-spawn.log 2>&1 &
643     spawned_pid=$!
644     disown "$spawned_pid" 2>/dev/null || true
645     sleep 3
646     if tailscaled_is_alive; then
647         tailscaled_up=1
648         echo "          tailscaled alive via direct spawn (pid=$spawned_pid)"
649         echo "          (NOTE: not under launchd supervision a reboot, re-run of"
650         echo "          this script, or ./toggle.sh is required to recreate."
651         echo "          Last 5 lines of stderr captured to /tmp/nullexit-tailscaled-spawn.log:)"
652         tail -n 5 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/~/ /' || true
653     fi
654 fi
655
656 if [ "$tailscaled_up" = 0 ]; then
657     echo "          HARD FAIL: tailscaled refuses to stay alive across every recovery tier"
658     echo "          launchctl-managed state (if available):"
659     launchctl print "gui/$UID/homebrew.mxcl.tailscale" 2>/dev/null \
660         | grep -E 'state =|last exit code =|runs =|path =' \
661         | sed 's/~/ /' \
662         || echo "          (no launchctl state available for homebrew.mxcl.tailscale)"
663     echo "          Last 10 lines of /tmp/nullexit-tailscaled-spawn.log (if any):"
664     tail -n 10 /tmp/nullexit-tailscaled-spawn.log 2>/dev/null | sed 's/~/ /' || echo "          (no log)"
665     echo "          Most likely remaining cause: macOS Local Network permission denied."

```

```

666     echo "          System Settings  Privacy & Security  Local Network  toggle Tailscale ON."
667     exit 7
668 fi
669 fi
670 sleep 3
671
672 # ----- 8. Verify -----
673 header "Step 8  Verify fresh install"
674 echo "    brew service tailscale: $(brew services list 2>/dev/null | awk '$1=="tailscale"{{print $2, $3}}' | tr
675 '\n' ' ' | sed 's/ $//')")
676 echo "    tailscaled process:      $(pgrep -lf tailscaled 2>/dev/null | head -1 || echo 'NONE (check brew logs)
677 ')")
678 ts_ver=""
679 if [ -x /opt/homebrew/bin/tailscale ]; then
680     ts_ver="$(/opt/homebrew/bin/tailscale version 2>/dev/null | head -1 | tr -d '\n')"
681 elif command -v tailscale >/dev/null 2>&1; then
682     ts_ver="$(tailscale version 2>/dev/null | head -1 | tr -d '\n')"
683 fi
684 echo "    tailscale version:      ${ts_ver:-binary missing}"
685 ts_status=""
686 if [ -x /opt/homebrew/bin/tailscale ]; then
687     ts_status="$(/opt/homebrew/bin/tailscale status 2>/dev/null | head -3)"
688 fi
689 if [ -n "$ts_status" ]; then
690     echo "    tailscale status:"
691     printf '%s\n' "$ts_status" | sed 's/~ / /'
692 else
693     echo "    tailscale status:      (binary missing or errored)"
694 fi
695
696 if { [ "$DRY" = 0 ] && /opt/homebrew/bin/tailscale status 2>/dev/null | grep -q '^Tailscale is stopped'; };
697 then
698     echo
699     echo "    NOTICE: 'tailscale status' still reports stopped after a fresh install."
700     echo "    Two equally likely causes:"
701     echo "    (a) macOS hasn't yet shown the first-run 'Local Network' prompt for"
702     echo "        Tailscale open System Settings Privacy & Security Local"
703     echo "        Network. Toggle Tailscale ON."
704     echo "    (b) macOS denied Local Network permission at some prior point."
705     echo "    Either way: System Settings Privacy & Security Local Network."
706     echo "    If the entry isn't present, invoking any tailscale command will"
707     echo "    re-trigger the prompt automatically."
708 fi
709
710 # ----- 9. Done -----
711 header "Done"
712 if [ "$DRY" = 1 ]; then
713     echo "    (dry-mode finished; no follow-up commands were issued.)"
714     exit 0
715 fi
716
717 cat <<'DONE'
718
719 Manual follow-ups (all of these; the script stops short of running sudo
720 tailscale up so you can verify the brew install is actually live first):
721
722 1. System Settings Privacy & Security Local Network Tailscale ON
723    macOS will normally pop the prompt the first time you run `tailscale`.
724    If it didn't, trigger it via the menu bar check.
725
726 2. Re-engage the tailnet + exit node:
727    sudo tailscale up --ssh=true --accept-dns=false \
728      --exit-node="$(cat .gateway_ip | tr -d '\r' | awk 'NR==1{{print $1; exit}}')" \
729      --exit-node-allow-lan-access=true
730
731 3. Re-run: ./toggle.sh
732    (re-installs the com.nullexit.wake-recovery LaunchAgent wiped
733    in Step 5 to clear the slate and re-engages the gateway)
734
735 4. Final verification:
736    curl -s https://www.cloudflare.com/cdn-cgi/trace | grep -E '^(warp|ip|colo)= '
737    expect: warp=on, ip=<Cloudflare IP>, colo=YYZ/YUL
738
739 Container-side (gateway / 100.100.21.8) Tailscale is unaffected by this
740 script. Other tailnet devices (S24, iPhone, Windows) are unaffected.
741
742 DONE

```

## 41 scripts/routing-fix.sh

```
1 #!/bin/sh
2 # routing-fix: Maintains routing for SOCKS5 proxy traffic through WARP (tun0)
3 #               and FORWARD rules for Tailscale exit-node return traffic.
4 #
5 # Runs inside the warp container's network namespace. This script ensures:
6 # 1. Table 200 has default via tun0 (for SOCKS5 proxy traffic from 172.18.0.2)
7 #    with the WARP endpoint exception via eth0 (prevents tunnel loop)
8 # 2. The CGNAT ip rule (100.64.0.0/10 table 52 tailscale routes) is maintained for Tailscale
9 # 3. FORWARD RELATED,ESTABLISHED rule allows return traffic for exit-node
10 #    forwarded connections (S24 tailscale0 eth0 host internet)
11 #
12 # CRITICAL: Does NOT touch the main table's default route. Gluetun handles that
13 #            internally via its own routing rules and WireGuard fwmark (0xca6c table 51820).
14 #            Overriding the main table default would create a loop: encrypted WireGuard
15 #            packets exiting tun0 would match default dev tun0 and re-enter the tunnel.
16
17 set -e
18
19 trap 'echo "routing-fix: Caught signal, exiting..."; exit 0' TERM INT QUIT
20
21 sleep 15 &
22 wait $! || true
23
24 echo "routing-fix: Setting up routing for SOCKS5 proxy through WARP (tun0)..."
25
26 # Dynamically detect Docker subnet from eth0
27 DOCKER_NET=$(ip route show dev eth0 | grep -v default | head -1 | awk '{print $1}')
28 DOCKER_GW=$(ip route show default | head -1 | awk '{print $3}')
29 WARP_ENDPOINT="${WARP_ENDPOINT}"
30 TAILSCALE_ALLOW_LAN_P2P="${TAILSCALE_ALLOW_LAN_P2P:-false}"
31 IP_BLOCKLIST_FILE="/userfilters/ip_blocklist.ipset"
32 LAST_IP_MTIME=""
33 echo "routing-fix: Docker network: ${DOCKER_NET}, gateway: ${DOCKER_GW}"
34
35 # Table 200: Used by ip rule 100 (from 172.18.0.2) for SOCKS5 proxy traffic
36 # WARP endpoint goes via eth0 to avoid tunnel loop, everything else via tun0
37 ip route add "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || \
38 ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
39 ip route replace default dev tun0 table 200 2>/dev/null || true
40
41 # Main table: Just ensure the Docker subnet route exists via eth0
42 # (gluetun owns the default route here do NOT touch it)
43 ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true
44
45 # FORWARD RELATED,ESTABLISHED: Allow return traffic for exit-node connections.
46 # The container has BOTH iptables (nftables backend) and iptables-legacy
47 # loaded. Gluetun's post-rules.txt uses the nftables backend, while Tailscale
48 # uses the legacy backend. Packets go through BOTH stacks, so we add the rule
49 # to both to ensure it survives regardless of which backend has policy DROP.
50 #
51 # Packet flow: S24 outgoing (tailscale0->eth0) ACCEPTed by first rule.
52 # Return traffic (eth0->eth0 via table 52 to host) needs RELATED,ESTABLISHED
53 # to match the conntrack entry created by the outgoing flow.
54 add_fwd_related_established() {
55     # nftables backend (used by gluetun's post-rules.txt)
56     if ! iptables -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
57         iptables -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
58     fi
59     # legacy backend (used by Tailscale's ts-forward)
60     if command -v iptables-legacy >/dev/null 2>&1; then
61         if ! iptables-legacy -C FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null; then
62             iptables-legacy -A FORWARD -m state --state RELATED,ESTABLISHED -j ACCEPT 2>/dev/null || true
63         fi
64     fi
65 }
66
67 add_fwd_related_established
68
69 # Policy Routing for Tailscale P2P:
70 # Tailscale sends its encrypted mesh UDP packets from port 41642.
71 # If these go out tun0 (WARP), Cloudflare's strict NAT drops the returning
72 # hole-punch packets, causing Tailscale to fail P2P and fall back to DERP relays (high latency).
73 # We mark packets originating from port 41642 and force them out the main table (eth0).
74 #
75 # LAN P2P is controlled by two sources (in priority order):
76 # 1. /app/.lan_p2p_detected written by watcher.sh on macOS on every network change
77 #    using system_profiler SPAirPortDataType (802.1x detection) + AP isolation probe. Overrides .env.
```

```

78 # 2. TAILSCALE_ALLOW_LAN_P2P env var (from .env) static fallback.
79 # When disabled, we explicitly DROP local RFC1918-destined Tailscale UDP so the
80 # phone cleanly falls back to DERP rather than getting poisoned by macOS gvpnproxy SNAT.
81 add_tailscale_p2p_bypass() {
82     # Dynamic override from watcher.sh (network auto-detection) takes priority over .env
83     local allow_p2p="${TAILSCALE_ALLOW_LAN_P2P:-false}"
84     if [ -f "/app/.lan_p2p_detected" ]; then
85         allow_p2p=$(cat "/app/.lan_p2p_detected" 2>/dev/null | tr -d '[:space:]')
86     fi
87
88     # Always allow P2P directly to the Docker host gateway (the Mac running this container).
89     # The gateway container is always co-located with the host Mac on the same machine,
90     # so direct Tailscale P2P via the Docker bridge is always safe and avoids DERP overhead.
91     # This RETURN rule must be inserted BEFORE the general DROP rules so it takes priority.
92     local default_gw
93     default_gw=$(ip route show default | awk '{print $3}')
94     if [ -n "$default_gw" ]; then
95         if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null; then
96             iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${default_gw}/32" -j RETURN 2>/dev/null || true
97         fi
98     fi
99
100     # Also explicitly whitelist all physical IPs of the host Mac (e.g. 172.17.52.223).
101     # The Tailscale control plane registers the Mac using its physical IP, so the
102     # container will attempt P2P handshakes using that physical IP instead of the
103     # Docker bridge IP. If this physical IP falls within 172.16.0.0/12, it would be dropped.
104     if [ -f "/app/.host_ips" ]; then
105         while IFS= read -r host_ip; do
106             [ -z "$host_ip" ] && continue
107             if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null; then
108                 iptables -t mangle -I OUTPUT 1 -p udp --sport 41642 -d "${host_ip}/32" -j RETURN 2>/dev/null || true
109             fi
110             done < "/app/.host_ips"
111     fi
112
113     if [ "${allow_p2p}" != "true" ]; then
114         # Drop Tailscale UDP packets destined for local subnets to prevent macOS gvpnproxy SNAT
115         # endpoint poisoning. When disabled (campus/enterprise WPA2-Enterprise with AP isolation),
116         # dropping these forces the S24 to use DERP relays reliably.
117         for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
118             if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; then
119                 iptables -t mangle -A OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null || true
120             fi
121             done
122         else
123             # Remove DROP rules if they exist (switching from restricted to allowed)
124             for subnet in "192.168.0.0/16" "10.0.0.0/8" "172.16.0.0/12"; do
125                 while iptables -t mangle -D OUTPUT -p udp --sport 41642 -d "$subnet" -j DROP 2>/dev/null; do ;; done
126             done
127         fi
128
129     # Bypass STUN and P2P packets to public IPs (DERP relays and remote peers) out eth0
130     if ! iptables -t mangle -C OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null; then
131         iptables -t mangle -A OUTPUT -p udp --sport 41642 -j MARK --set-mark 0x8888 2>/dev/null || true
132     fi
133     if ! ip rule show | grep -q 'fwmark 0x8888'; then
134         ip rule add fwmark 0x8888 lookup 254 pref 98 2>/dev/null || true
135     fi
136 }
137
138 add_tailscale_p2p_bypass
139
140 add_onion_transparent_proxy() {
141     for ipt in iptables iptables-legacy; do
142         command -v "$ipt" >/dev/null 2>&1 || continue
143         if ! $ipt -t nat -C PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null; then
144             $ipt -t nat -I PREROUTING -d 198.18.0.0/15 -p tcp -j REDIRECT --to-ports 9040 2>/dev/null || true
145         fi
146         if ! $ipt -C FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null; then
147             $ipt -I FORWARD -d 198.18.0.0/15 -j ACCEPT 2>/dev/null || true
148         fi
149     done
150 }
151
152 add_tailscale_p2p_bypass
153 add_onion_transparent_proxy
154
155 create_log_drop_chain() {
156     local ipt="$1"
157     local chain="$2"
158     local prefix="$3"

```

```

159
160 if ! $ipt -N "$chain" 2>/dev/null; then
161     $ipt -F "$chain" 2>/dev/null || true
162 fi
163 $ipt -A "$chain" -j LOG --log-prefix "$prefix"
164 $ipt -A "$chain" -j DROP 2>/dev/null || true
165 }
166
167 add_country_block() {
168     # To add a country to the blocklist, simply define the 'BLOCKED_COUNTRIES' variable in .env with 2-letter ISO
169     # country codes.
170     # You can find the country code by looking at the filename on http://www.ipdeny.com/ipblocks/data/countries/
171     # For example, to block China (cn) and Russia (ru), set: BLOCKED_COUNTRIES="cn ru" in your .env file
172     for zone in $BLOCKED_COUNTRIES; do
173         set_name="block_${zone}"
174         zone_upper=$(echo "$zone" | tr 'a-z' 'A-Z')
175         chain_name="LOG_DROP_${zone_upper}"
176         log_prefix="IP_BLOCK_${zone_upper}: "
177
178         # Create the ipset if it doesn't exist
179         if ! ipset list "$set_name" >/dev/null 2>&1; then
180             ipset create "$set_name" hash:net 2>/dev/null || true
181             echo "routing-fix: Downloading ${zone_upper} IPs for blocklist..."
182             curl -sS "http://www.ipdeny.com/ipblocks/data/countries/${zone}.zone" | grep -v '^#' | while read -r
183             subnet; do
184                 [ -n "$subnet" ] && ipset add "$set_name" "$subnet" 2>/dev/null || true
185             done
186         fi
187
188         # Apply LOG_DROP rule to both backends
189         for ipt in iptables iptables-legacy; do
190             command -v "$ipt" >/dev/null 2>&1 || continue
191
192             create_log_drop_chain "$ipt" "$chain_name" "$log_prefix"
193
194             # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
195             while $ipt -D FORWARD -m set --match-set "$set_name" dst -j DROP 2>/dev/null; do ;; done
196
197             if ! $ipt -C FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null; then
198                 $ipt -I FORWARD -m set --match-set "$set_name" dst -j "$chain_name" 2>/dev/null || true
199             fi
200         done
201     done
202 }
203
204 add_ip_blocklist() {
205     # Only reload when the compiled file has actually changed (mtime check).
206     # This prevents a pointless ipset restore on every 30-second loop tick.
207     if [ ! -f "$IP_BLOCKLIST_FILE" ]; then
208         return 0
209     fi
210
211     CURRENT_MTIME=$(stat -c %Y "$IP_BLOCKLIST_FILE" 2>/dev/null || echo "0")
212     if [ "$CURRENT_MTIME" != "$LAST_IP_MTIME" ]; then
213         echo "routing-fix: IP blocklist changed, reloading..."
214         # ipset restore handles the atomic swap internally:
215         # create_new populate swap with live destroy_new
216         if ipset restore < "$IP_BLOCKLIST_FILE" 2>/dev/null; then
217             LAST_IP_MTIME="$CURRENT_MTIME"
218             echo "routing-fix: IP blocklist loaded ($(ipset list block_malicious | grep -c '^[0-9]') entries)."
219         else
220             echo "routing-fix: Warning: ipset restore failed. Retrying next cycle."
221             return 1
222         fi
223     fi
224
225     # Apply FORWARD rules in BOTH iptables backends.
226     for ipt in iptables iptables-legacy; do
227         command -v "$ipt" >/dev/null 2>&1 || continue
228
229         create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_DST "IP_BLOCK_MALICIOUS_DST: "
230         create_log_drop_chain "$ipt" LOG_DROP_MALICIOUS_SRC "IP_BLOCK_MALICIOUS_SRC: "
231
232         # Clean up old plain DROP rules to ensure we hit the log-and-drop chains
233         while $ipt -D FORWARD -m set --match-set block_malicious dst -j DROP 2>/dev/null; do ;; done
234         while $ipt -D FORWARD -m set --match-set block_malicious src -j DROP 2>/dev/null; do ;; done
235
236         if ! $ipt -C FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null; then
237             $ipt -I FORWARD -m set --match-set block_malicious dst -j LOG_DROP_MALICIOUS_DST 2>/dev/null || true
238         fi
239     done
240 }

```



```

238     if ! $ipt -C FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null; then
239         $ipt -I FORWARD -m set --match-set block_malicious src -j LOG_DROP_MALICIOUS_SRC 2>/dev/null || true
240     fi
241 done
242 }
243
244 # Start background block logger (DNS + IP)
245 if ! pgrep -f "nullexit-logger" >/dev/null; then
246     echo "routing-fix: Starting background block logger..."
247     nullexit-logger &
248 fi
249
250 add_country_block
251 add_ip_blocklist
252
253 echo 'routing-fix: Routes applied.'
254
255 UNHEALTHY_COUNT=0
256
257 # Re-assert loop (every 30 seconds)
258 while true; do
259     # Table 200 routes
260     ip route replace "${WARP_ENDPOINT}" via "${DOCKER_GW}" dev eth0 table 200 2>/dev/null || true
261
262     # Tunnel health check
263     if ! curl -m 3 -sS --interface tun0 https://cloudflare.com/cdn-cgi/trace >/dev/null 2>&1; then
264         UNHEALTHY_COUNT=$((UNHEALTHY_COUNT + 1))
265     else
266         UNHEALTHY_COUNT=0
267         if [ -f "/app/TUNNEL_FAILED_CLOSED.marker" ]; then
268             rm -f "/app/TUNNEL_FAILED_CLOSED.marker"
269         fi
270     fi
271
272     if [ "$UNHEALTHY_COUNT" -ge 3 ]; then
273         echo "routing-fix: Tunnel health check failed $UNHEALTHY_COUNT times. Failing closed."
274         ip route del default dev tun0 table 200 2>/dev/null || true
275         ip route replace blackhole default table 200 2>/dev/null || true
276         touch "/app/TUNNEL_FAILED_CLOSED.marker"
277     else
278         ip route replace default dev tun0 table 200 2>/dev/null || true
279     fi
280
281     # Main table: keep Docker subnet route
282     ip route replace "${DOCKER_NET}" dev eth0 2>/dev/null || true
283
284     # CGNAT ip rule for Tailscale
285     if ! ip rule show | grep -q '100.64.0.0/10'; then
286         ip rule add to 100.64.0.0/10 lookup 52 pref 99 2>/dev/null || true
287         echo 'routing-fix: CGNAT rule missing, re-injected'
288     fi
289
290     # FORWARD RELATED, ESTABLISHED: Allow return traffic in both iptables
291     # backends. Gluetun may reset the nftables ruleset on VPN reconnect,
292     # and Tailscale may reset the legacy ruleset on auth refresh.
293     add_fwd_related_established
294
295     # Ensure P2P packets bypass WARP to maintain low latency
296     add_tailscale_p2p_bypass
297
298     # Ensure transparent proxying for .onion is active
299     add_onion_transparent_proxy
300
301     # Enforce country blocklist
302     add_country_block
303
304     # Enforce IP blocklist (reloads automatically when file changes)
305     add_ip_blocklist
306
307     sleep 30 &
308     wait $! || true
309 done

```

## 42 scripts/rule-compiler/go.mod

```
1 module nullexit/rule-compiler
```

```
2
3 go 1.22
```

## 43 scripts/rule-compiler/main.go

```
1 package main
2
3 import (
4     "bufio"
5     "bytes"
6     "crypto/md5"
7     "encoding/base64"
8     "encoding/hex"
9     "fmt"
10    "io"
11    "net/http"
12    "net/netip"
13    "os"
14    "os/exec"
15    "path/filepath"
16    "regexp"
17    "sort"
18    "strings"
19    "sync"
20    "time"
21 )
22
23 const (
24     BlackListPath = "black_list.txt"
25     WhiteListPath = "white_list.txt"
26     OutputPath    = "adguard/work/userfilters/compiled_rules.txt"
27     IpOutputPath  = "adguard/work/userfilters/ip_blocklist.ipset"
28     IpCacheDir    = "adguard/work/userfilters/cache/ip"
29     CacheDir      = "adguard/work/userfilters/cache"
30 )
31
32 var CoreLists = []string{
33     "https://raw.githubusercontent.com/anudeepND/blacklist/master/adservers.txt",
34     "https://raw.githubusercontent.com/anudeepND/blacklist/master/facebook.txt",
35     "https://raw.githubusercontent.com/crazy-max/WindowsSpyBlocker/master/data/hosts/spy.txt",
36     "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.samsung.txt",
37     "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/adblock/native.apple.txt",
38     "https://abp.oisd.nl/basic/",
39     "https://adaway.org/hosts.txt",
40     "https://pgl.yoyo.org/adserver/serverlist.php?hostformat=adblockplus&showintro=1&mimetype=plaintext",
41     "https://raw.githubusercontent.com/nextdns/cname-cloaking-blocklist/master/domains",
42 }
43
44 var MediumAdditions = []string{
45     "https://raw.githubusercontent.com/lightswitch05/hosts/master/docs/lists/facebook-extended.txt",
46     "https://someonewhocares.org/hosts/zero/hosts",
47     "https://urlhaus.abuse.ch/downloads/hostfile/",
48 }
49
50 var HeavyAdditions = []string{
51     "https://raw.githubusercontent.com/StevenBlack/hosts/master/hosts",
52     "https://raw.githubusercontent.com/Perflyst/PiHoleBlocklist/master/SmartTV-AGH.txt",
53     "https://raw.githubusercontent.com/DandelionSprout/adfilt/master/GameConsoleAdblockList.txt",
54 }
55
56 var IpBlockLists = []string{
57     "https://feodotracker.abuse.ch/downloads/ipblocklist.txt",
58     "https://www.spamhaus.org/drop/drop.txt",
59     "https://rules.emergingthreats.net/fwrules/emerging-Block-IPs.txt",
60     "https://cinsscore.com/list/ci-badguys.txt",
61 }
62
63 func getProfiles() map[string][]string {
64     profiles := make(map[string][]string)
65
66     light := append([]string(nil), CoreLists...)
67     light = append(light, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/light.txt")
68     profiles["light"] = light
69
70     medium := append([]string(nil), CoreLists...)
71     medium = append(medium, MediumAdditions...)
```

```

72 medium = append(medium, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/multi.txt")
73 profiles["medium"] = medium
74
75 heavy := append([]string(nil), CoreLists...)
76 heavy = append(heavy, MediumAdditions...)
77 heavy = append(heavy, HeavyAdditions...)
78 heavy = append(heavy, "https://raw.githubusercontent.com/hagezi/dns-blocklists/main/domains/pro.txt")
79 profiles["heavy"] = heavy
80
81 return profiles
82 }
83
84 func loadEnvProfile() string {
85     profile := "medium"
86
87     if bytesData, err := os.ReadFile(".env"); err == nil {
88         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
89         for scanner.Scan() {
90             line := strings.TrimSpace(scanner.Text())
91             if line != "" && !strings.HasPrefix(line, "#") {
92                 parts := strings.SplitN(line, "=", 2)
93                 if len(parts) == 2 {
94                     key := strings.TrimSpace(parts[0])
95                     val := strings.TrimSpace(parts[1])
96                     if key == "GATEWAY_RULE_PROFILE" {
97                         val = strings.Trim(val, "\"")
98                         profile = strings.ToLower(val)
99                     }
100                 }
101             }
102         }
103     } else if !os.IsNotExist(err) {
104         fmt.Printf("Warning: Failed to read .env file for profile configuration (%v). Using default.\n", err)
105     }
106
107     if envVal := os.Getenv("GATEWAY_RULE_PROFILE"); envVal != "" {
108         profile = strings.ToLower(envVal)
109     }
110
111     profiles := getProfiles()
112     if _, exists := profiles[profile]; !exists {
113         fmt.Printf("Warning: Profile '%s' is invalid. Falling back to 'medium'.\n", profile)
114         profile = "medium"
115     }
116     return profile
117 }
118
119 func loadDomains(filepath string) map[string]bool {
120     domains := make(map[string]bool)
121     if bytesData, err := os.ReadFile(filepath); err == nil {
122         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
123         for scanner.Scan() {
124             line := strings.TrimSpace(scanner.Text())
125             if line != "" && !strings.HasPrefix(line, "#") {
126                 domains[strings.ToLower(line)] = true
127             }
128         }
129     }
130     return domains
131 }
132
133 func parseDomainsFromContent(content string) map[string]bool {
134     domains := make(map[string]bool)
135     scanner := bufio.NewScanner(strings.NewReader(content))
136     ipRegex := regexp.MustCompile(`^\d{1,3}\.\d{1,3}\.\d{1,3}\.\d{1,3}$`)
137     agRegex := regexp.MustCompile(`^\|([a-z0-9.-]+)\|` )
138     domainRegex := regexp.MustCompile(`^[a-z0-9.-]+$`)
139
140     for scanner.Scan() {
141         line := strings.TrimSpace(scanner.Text())
142         if line == "" || strings.HasPrefix(line, "!") || strings.HasPrefix(line, "[") {
143             continue
144         }
145         parts := strings.SplitN(line, "#", 2)
146         line = strings.TrimSpace(parts[0])
147         if line == "" {
148             continue
149         }
150
151         fields := strings.Fields(line)
152         var rule string

```

```

153     if len(fields) >= 2 {
154         if ipRegex.MatchString(fields[0]) {
155             rule = fields[1]
156         } else {
157             rule = fields[0]
158         }
159     } else {
160         rule = fields[0]
161     }
162
163     rule = strings.TrimSpace(strings.ToLower(rule))
164
165     var domain string
166     if m := agRegex.FindStringSubmatch(rule); m != nil {
167         domain = m[1]
168     } else if domainRegex.MatchString(rule) {
169         domain = rule
170     } else {
171         domain = rule
172     }
173
174     if domain != "" && domain != "localhost" && domain != "0.0.0.0" && domain != "127.0.0.1" && domain != "
175         broadcasthost" {
176         domains[domain] = true
177     }
178     return domains
179 }
180
181 // WARNING: Parsing AdGuardHome.yaml using regex instead of a proper YAML parser
182 // is brittle and assumes the exact format currently emitted by AdGuard.
183 // If an AdGuard version bump updates the configuration schema format,
184 // this parser will fail silently or loudly and should be refactored to use gopkg.in/yaml.v3.
185 func getAdguardNativeLists() []string {
186     yamlPath := "adguard/conf/AdGuardHome.yaml"
187     var enabledUrls []string
188     if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
189         return enabledUrls
190     }
191
192     contentBytes, err := os.ReadFile(yamlPath)
193     if err != nil {
194         fmt.Printf("Warning: Failed to parse AdGuardHome.yaml for native lists (%v)\n", err)
195         return enabledUrls
196     }
197     content := string(contentBytes)
198
199     filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*)"?(?![a-zA-Z_]+\+|\z)`
200     match := filtersRegex.FindStringSubmatch(content)
201     if len(match) > 1 {
202         filtersText := match[1]
203         blocks := strings.Split(filtersText, "\n - ")
204         urlRegex := regexp.MustCompile(`url:\s*(\S+)`)
205
206         for _, block := range blocks {
207             if strings.TrimSpace(block) == "" {
208                 continue
209             }
210             if strings.Contains(block, "enabled: true") {
211                 urlMatch := urlRegex.FindStringSubmatch(block)
212                 if len(urlMatch) > 1 {
213                     url := urlMatch[1]
214                     if !strings.Contains(url, "compiled_rules.txt") {
215                         enabledUrls = append(enabledUrls, url)
216                     }
217                 }
218             }
219         }
220     }
221     return enabledUrls
222 }
223
224 func getAdguardCompiledRulesCachePath() string {
225     yamlPath := "adguard/conf/AdGuardHome.yaml"
226     if _, err := os.Stat(yamlPath); os.IsNotExist(err) {
227         return ""
228     }
229
230     contentBytes, err := os.ReadFile(yamlPath)
231     if err != nil {
232         fmt.Printf("Warning: Failed to resolve compiled_rules.txt cache path (%v)\n", err)

```

```

233     return ""
234 }
235 content := string(contentBytes)
236
237 filtersRegex := regexp.MustCompile(`(?ms)^filters:(.*?)(?:^[a-zA-Z_]+:|\z)` )
238 match := filtersRegex.FindStringSubmatch(content)
239 if len(match) > 1 {
240     blocks := strings.Split(match[1], "\n - ")
241     idRegex := regexp.MustCompile(`id:\s*(\d+)`)
242     for _, block := range blocks {
243         if !strings.Contains(block, "compiled_rules.txt") {
244             continue
245         }
246         idMatch := idRegex.FindStringSubmatch(block)
247         if len(idMatch) > 1 {
248             return fmt.Sprintf("adguard/work/data/filters/%s.txt", idMatch[1])
249         }
250     }
251 }
252 return ""
253 }
254
255 func handleFetchError(urlStr string, cacheFile string, err error) map[string]bool {
256     if _, statErr := os.Stat(cacheFile); statErr == nil {
257         fmt.Printf(" -> Warning: Failed to fetch %s (%v). Falling back to expired local cache.\n", urlStr, err)
258         content, readErr := os.ReadFile(cacheFile)
259         if readErr == nil {
260             domains := parseDomainsFromContent(string(content))
261             fmt.Printf(" -> Loaded from expired cache (%d domains).\n", len(domains))
262             return domains
263         }
264         fmt.Printf(" -> Warning: Failed to read expired cache (%v). Using local lists only.\n", readErr)
265     } else {
266         fmt.Printf(" -> Warning: Failed to fetch %s (%v). Using local lists only for this source.\n", urlStr, err)
267     }
268     return make(map[string]bool)
269 }
270
271 func fetchRemoteDomains(urlStr string) map[string]bool {
272     os.MkdirAll(CacheDir, 0755)
273
274     hash := md5.Sum([]byte(urlStr))
275     urlHash := hex.EncodeToString(hash[:])
276     cacheFile := filepath.Join(CacheDir, fmt.Sprintf("%s.txt", urlHash))
277
278     if stat, err := os.Stat(cacheFile); err == nil {
279         fileAge := time.Since(stat.ModTime()).Seconds()
280         if fileAge < 86400 {
281             fmt.Printf("Loading remote blacklist from cache: %s\n", urlStr)
282             content, err := os.ReadFile(cacheFile)
283             if err == nil {
284                 domains := parseDomainsFromContent(string(content))
285                 fmt.Printf(" -> Loaded from cache (copy is %.1f hours old, %d domains).\n", fileAge/3600.0, len(domains))
286             }
287             return domains
288         }
289         fmt.Printf(" -> Warning: Failed to read cache file %s (%v). Re-fetching...\n", cacheFile, err)
290     }
291
292     fmt.Printf("Fetching remote blacklist from: %s ...\n", urlStr)
293
294     req, err := http.NewRequest("GET", urlStr, nil)
295     if err != nil {
296         return handleFetchError(urlStr, cacheFile, err)
297     }
298     req.Header.Set("User-Agent", "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_7)")
299
300     client := &http.Client{Timeout: 15 * time.Second}
301     resp, err := client.Do(req)
302     if err != nil {
303         return handleFetchError(urlStr, cacheFile, err)
304     }
305     defer resp.Body.Close()
306
307     if resp.StatusCode != 200 {
308         return handleFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
309     }
310
311     bodyBytes, err := io.ReadAll(resp.Body)
312     if err != nil {

```

```

313     return handleFetchError(urlStr, cacheFile, err)
314 }
315
316 content := string(bodyBytes)
317 domains := parseDomainsFromContent(content)
318
319 if len(domains) < 10 {
320     return handleFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only found %d domains. Possible
321         404 or bad URL", len(domains)))
322 }
323
324 err = os.WriteFile(cacheFile, bodyBytes, 0644)
325 if err != nil {
326     fmt.Printf("-> Warning: Failed to save cache file (%v)\n", err)
327 }
328
329 fmt.Printf("-> Successfully fetched and cached %d domains.\n", len(domains))
330 return domains
331 }
332
333 func optimizeSubdomains(domains map[string]bool, listName string) map[string]bool {
334     fmt.Printf("Optimizing %s by removing redundant subdomains...\n", listName)
335     rawCount := len(domains)
336     if rawCount == 0 {
337         return make(map[string]bool)
338     }
339
340     optimized := make(map[string]bool)
341     for domain := range domains {
342         if strings.ContainsAny(domain, "|~@$/%") {
343             optimized[domain] = true
344             continue
345         }
346
347         parts := strings.Split(domain, ".")
348         hasParent := false
349
350         for i := 1; i < len(parts)-1; i++ {
351             parent := strings.Join(parts[i:], ".")
352             if domains[parent] {
353                 hasParent = true
354                 break
355             }
356         }
357
358         if !hasParent {
359             optimized[domain] = true
360         }
361     }
362
363     saved := rawCount - len(optimized)
364     reduction := float64(0)
365     if rawCount > 0 {
366         reduction = (float64(saved) / float64(rawCount)) * 100
367     }
368     fmt.Printf("-> Reduced %s from %d to %d domains (-%d / %.1f%% reduction).\n", listName, rawCount, len(
369         optimized), saved, reduction)
370     return optimized
371 }
372
373 func parseIpsFromContent(content string) map[string]bool {
374     ips := make(map[string]bool)
375     scanner := bufio.NewScanner(strings.NewReader(content))
376
377     for scanner.Scan() {
378         line := strings.TrimSpace(scanner.Text())
379         if line == "" || strings.HasPrefix(line, "#") || strings.HasPrefix(line, ";") {
380             continue
381         }
382
383         fields := strings.Fields(line)
384         if len(fields) > 0 {
385             token := strings.TrimRight(fields[0], ";,")
386             if _, err := netip.ParsePrefix(token); err == nil {
387                 ips[token] = true
388             } else if _, err := netip.ParseAddr(token); err == nil {
389                 ips[token] = true
390             }
391         }
392     }
393     return ips

```

```

392 }
393
394 func handleIpFetchError(urlStr string, cacheFile string, err error) map[string]bool {
395     if _, statErr := os.Stat(cacheFile); statErr == nil {
396         fmt.Printf("-> Warning: fetch failed (%v). Falling back to stale cache.\n", err)
397         content, readErr := os.ReadFile(cacheFile)
398         if readErr == nil {
399             ips := parseIpsFromContent(string(content))
400             fmt.Printf("-> %d IPs loaded from stale cache.\n", len(ips))
401             return ips
402         }
403         fmt.Printf("-> Warning: stale cache also failed (%v).\n", readErr)
404     } else {
405         fmt.Printf("-> Warning: %s failed (%v). No cache available. Skipping.\n", urlStr, err)
406     }
407     return make(map[string]bool)
408 }
409
410 func fetchRemoteIps(urlStr string) map[string]bool {
411     os.MkdirAll(IpCacheDir, 0755)
412
413     hash := md5.Sum([]byte(urlStr))
414     urlHash := hex.EncodeToString(hash[:])
415     cacheFile := filepath.Join(IpCacheDir, fmt.Sprintf("%s.txt", urlHash))
416
417     if stat, err := os.Stat(cacheFile); err == nil {
418         fileAge := time.Since(stat.ModTime()).Seconds()
419         if fileAge < 86400 {
420             fmt.Printf("Loading IP feed from cache: %s\n", urlStr)
421             content, err := os.ReadFile(cacheFile)
422             if err == nil {
423                 ips := parseIpsFromContent(string(content))
424                 fmt.Printf("-> %d IPs (cache is %.1fh old).\n", len(ips), fileAge/3600.0)
425                 return ips
426             }
427             fmt.Printf("-> Warning: cache read failed (%v). Re-fetching...\n", err)
428         }
429     }
430
431     fmt.Printf("Fetching IP feed: %s ...\n", urlStr)
432     req, err := http.NewRequest("GET", urlStr, nil)
433     if err != nil {
434         return handleIpFetchError(urlStr, cacheFile, err)
435     }
436     req.Header.Set("User-Agent", "Mozilla/5.0")
437
438     client := &http.Client{Timeout: 15 * time.Second}
439     resp, err := client.Do(req)
440     if err != nil {
441         return handleIpFetchError(urlStr, cacheFile, err)
442     }
443     defer resp.Body.Close()
444
445     if resp.StatusCode != 200 {
446         return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("HTTP %d", resp.StatusCode))
447     }
448
449     bodyBytes, err := io.ReadAll(resp.Body)
450     if err != nil {
451         return handleIpFetchError(urlStr, cacheFile, err)
452     }
453
454     content := string(bodyBytes)
455     ips := parseIpsFromContent(content)
456     if len(ips) < 1 {
457         return handleIpFetchError(urlStr, cacheFile, fmt.Errorf("Sanity check failed: only %d IPs found. Possible 404", len(ips)))
458     }
459
460     err = os.WriteFile(cacheFile, bodyBytes, 0644)
461     if err != nil {
462         fmt.Printf("-> Warning: cache write failed (%v)\n", err)
463     }
464
465     fmt.Printf("-> %d IPs fetched and cached.\n", len(ips))
466     return ips
467 }
468
469 func compileIpBlocklist() int {
470     fmt.Println("\n IP Blocklist Compilation ")
471

```



```

472 allIps := make(map[string]bool)
473 var mu sync.Mutex
474 var wg sync.WaitGroup
475
476 maxThreads := 8
477 if len(IpBlockLists) < 8 {
478     maxThreads = len(IpBlockLists)
479 }
480 semaphore := make(chan struct{}, maxThreads)
481
482 for _, urlStr := range IpBlockLists {
483     wg.Add(1)
484     go func(u string) {
485         defer wg.Done()
486         semaphore <- struct{}{}
487         defer func() { <-semaphore }()
488
489         ips := fetchRemoteIps(u)
490         mu.Lock()
491         for ip := range ips {
492             allIps[ip] = true
493         }
494         mu.Unlock()
495     }(urlStr)
496 }
497 wg.Wait()
498
499 fmt.Printf("Total unique entries before filtering: %d\n", len(allIps))
500
501 reservedStr := []string{
502     "10.0.0.0/8",
503     "172.16.0.0/12",
504     "192.168.0.0/16",
505     "127.0.0.0/8",
506     "169.254.0.0/16",
507     "100.64.0.0/10",
508     "0.0.0.0/8",
509 }
510 var reserved []netip.Prefix
511 for _, s := range reservedStr {
512     p, _ := netip.ParsePrefix(s)
513     reserved = append(reserved, p)
514 }
515
516 cleanIps := make(map[string]bool)
517 for entry := range allIps {
518     var prefix netip.Prefix
519     p, err := netip.ParsePrefix(entry)
520     if err != nil {
521         addr, err2 := netip.ParseAddr(entry)
522         if err2 != nil {
523             continue
524         }
525         prefix = netip.PrefixFrom(addr, addr.BitLen())
526     } else {
527         prefix = p
528     }
529
530     overlaps := false
531     for _, r := range reserved {
532         if r.Overlaps(prefix) {
533             overlaps = true
534             break
535         }
536     }
537     if !overlaps {
538         cleanIps[entry] = true
539     }
540 }
541
542 removed := len(allIps) - len(cleanIps)
543 if removed > 0 {
544     fmt.Printf(" -> Removed %d private/reserved entries.\n", removed)
545 }
546 fmt.Printf(" -> Final IP blocklist: %d entries.\n", len(cleanIps))
547
548 os.MkdirAll(filepath.Dir(IpOutputPath), 0755)
549
550 f, err := os.Create(IpOutputPath)
551 if err != nil {
552     fmt.Printf("Error writing IP output file: %v\n", err)

```

```

553     return len(cleanIps)
554 }
555 defer f.Close()
556
557 f.WriteString("# nullexit Compiled IP Blocklist\n")
558 f.WriteString("# Sources: Feodo Tracker, Spamhaus DROP/EDROP, Emerging Threats, CINS\n")
559 f.WriteString(fmt.Sprintf("# Entries: %d\n", len(cleanIps)))
560
561 f.WriteString("create block_malicious_new hash:net maxelem 200000 -exist\n")
562
563 var sortedIps []string
564 for ip := range cleanIps {
565     sortedIps = append(sortedIps, ip)
566 }
567 sort.Strings(sortedIps)
568
569 for _, ip := range sortedIps {
570     f.WriteString(fmt.Sprintf("add block_malicious_new %s -exist\n", ip))
571 }
572
573 f.WriteString("create block_malicious hash:net maxelem 200000 -exist\n")
574 f.WriteString("swap block_malicious block_malicious_new\n")
575 f.WriteString("destroy block_malicious_new\n")
576
577 fmt.Printf("IP blocklist written to %s\n", IpOutputPath)
578 return len(cleanIps)
579 }
580
581 func main() {
582     profile := loadEnvProfile()
583     fmt.Printf("Active memory profile: '%s'\n", strings.ToUpper(profile))
584
585     localBlackList := loadDomains(BlackListPath)
586     whiteList := loadDomains(WhiteListPath)
587
588     blackList := make(map[string]bool)
589     for k := range localBlackList {
590         blackList[k] = true
591     }
592
593     profiles := getProfiles()
594     urlsToFetch := profiles[profile]
595
596     fmt.Printf("\nStarting concurrent downloads for %d lists...\n", len(urlsToFetch))
597     startTime := time.Now()
598
599     var mu sync.Mutex
600     var wg sync.WaitGroup
601
602     maxThreads := 16
603     if len(urlsToFetch) < 16 {
604         maxThreads = len(urlsToFetch)
605     }
606     semaphore := make(chan struct{}, maxThreads)
607
608     for _, urlStr := range urlsToFetch {
609         wg.Add(1)
610         go func(u string) {
611             defer wg.Done()
612             semaphore <- struct{}{}
613             defer func() { <-semaphore }()
614
615             domains := fetchRemoteDomains(u)
616             mu.Lock()
617             for d := range domains {
618                 blackList[d] = true
619             }
620             mu.Unlock()
621         }(urlStr)
622     }
623     wg.Wait()
624
625     fmt.Printf("Finished concurrent downloads in %.2f seconds using %d threads.\n", time.Since(startTime).Seconds(), maxThreads)
626
627     adguardNativeUrls := getAdguardNativeLists()
628     var adguardNativeDomains map[string]bool
629     if len(adguardNativeUrls) > 0 {
630         fmt.Printf("\nFetching %d AdGuard native list(s) for deduplication...\n", len(adguardNativeUrls))
631         adguardNativeDomains = make(map[string]bool)
632     }

```

```

633 maxNativeThreads := 4
634 if len(adguardNativeUrls) < 4 {
635     maxNativeThreads = len(adguardNativeUrls)
636 }
637 nativeSemaphore := make(chan struct{}, maxNativeThreads)
638
639 for _, urlStr := range adguardNativeUrls {
640     wg.Add(1)
641     go func(u string) {
642         defer wg.Done()
643         nativeSemaphore <- struct{}{}
644         defer func() { <-nativeSemaphore }()
645
646         domains := fetchRemoteDomains(u)
647         mu.Lock()
648         for d := range domains {
649             adguardNativeDomains[d] = true
650         }
651         mu.Unlock()
652     }(urlStr)
653 }
654 wg.Wait()
655
656 if len(adguardNativeDomains) > 0 {
657     fmt.Println("Deduplicating compiled rules against AdGuard native lists...")
658     originalSize := len(blackList)
659     for d := range adguardNativeDomains {
660         delete(blackList, d)
661     }
662     fmt.Printf(" -> Removed %d redundant rules already covered by AdGuard.\n", originalSize-len(blackList))
663 }
664 }
665
666 var contradictions []string
667 for d := range whiteList {
668     if blackList[d] {
669         contradictions = append(contradictions, d)
670     }
671 }
672
673 if len(contradictions) > 0 {
674     fmt.Printf("Detected %d contradictions. Whitelist taking priority.\n", len(contradictions))
675     sort.Strings(contradictions)
676     for _, domain := range contradictions {
677         if localBlackList[domain] {
678             fmt.Printf(" -> Removed local blacklist domain: %s\n", domain)
679         }
680         delete(blackList, domain)
681     }
682 }
683
684 blackList = optimizeSubdomains(blackList, "blacklist")
685 whiteList = optimizeSubdomains(whiteList, "whitelist")
686
687 os.MkdirAll(filepath.Dir(OutputPath), 0755)
688 outFile, err := os.Create(OutputPath)
689 if err != nil {
690     fmt.Printf("Error creating output file: %v\n", err)
691     return
692 }
693 defer outFile.Close()
694
695 outFile.WriteString("! Custom Compiled Rules (Auto-Generated)\n")
696 outFile.WriteString(fmt.Sprintf("! Memory Profile: %s\n", strings.ToUpper(profile)))
697 outFile.WriteString(fmt.Sprintf("! Total Block Rules: %d\n", len(blackList)))
698 outFile.WriteString(fmt.Sprintf("! Native AdGuard Rules: %d\n", len(adguardNativeDomains)))
699 outFile.WriteString(fmt.Sprintf("! Total Whitelist Rules: %d\n", len(whiteList)))
700
701 outFile.WriteString("! --- Blacklist Rules ---\n")
702
703 var sortedBlacklist []string
704 for d := range blackList {
705     sortedBlacklist = append(sortedBlacklist, d)
706 }
707 sort.Strings(sortedBlacklist)
708
709 for _, domain := range sortedBlacklist {
710     if strings.Contains(domain, "$") {
711         parts := strings.SplitN(domain, "$", 2)
712         dom, mod := parts[0], parts[1]
713         if strings.HasPrefix(dom, "||") {

```

```

714         if strings.HasSuffix(dom, "~") {
715             outFile.WriteString(fmt.Sprintf("%s%s\n", dom, mod))
716         } else {
717             outFile.WriteString(fmt.Sprintf("%s~%s\n", dom, mod))
718         }
719     } else {
720         outFile.WriteString(fmt.Sprintf("||%s~%s\n", dom, mod))
721     }
722 } else if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|")
723 ) || strings.HasSuffix(domain, "~") {
724     outFile.WriteString(fmt.Sprintf("%s\n", domain))
725 } else {
726     outFile.WriteString(fmt.Sprintf("||%s~\n", domain))
727 }
728 }
729 outFile.WriteString("\n! --- Whitelist Rules ---\n")
730 var sortedWhitelist []string
731 for d := range whitelist {
732     sortedWhitelist = append(sortedWhitelist, d)
733 }
734 sort.Strings(sortedWhitelist)
735
736 for _, domain := range sortedWhitelist {
737     if strings.HasPrefix(domain, "/") || strings.HasPrefix(domain, "|") || strings.HasSuffix(domain, "|") ||
738 strings.HasSuffix(domain, "~") || strings.HasPrefix(domain, "@@") {
739         rule := domain
740         if !strings.HasPrefix(domain, "@@") {
741             rule = "@@" + domain
742         }
743         outFile.WriteString(fmt.Sprintf("%s\n", rule))
744     } else {
745         outFile.WriteString(fmt.Sprintf("@@||%s~\n", domain))
746     }
747 }
748
749 fmt.Printf("\nSuccessfully compiled %d block rules and %d allow rules to %s\n", len(blackList), len(whitelist),
750 OutputPath)
751
752 cachedFilterPath := getAdguardCompiledRulesCachePath()
753 if cachedFilterPath != "" {
754     input, err := os.ReadFile(OutputPath)
755     if err == nil {
756         err = os.WriteFile(cachedFilterPath, input, 0644)
757         if err == nil {
758             fmt.Printf("Updated AdGuard filter cache: %s\n", cachedFilterPath)
759         } else {
760             fmt.Printf("Warning: Could not update AdGuard filter cache %s (%v)\n", cachedFilterPath, err)
761         }
762     } else {
763         fmt.Printf("Warning: Could not read OutputPath to copy to cache (%v)\n", err)
764     }
765 } else {
766     fmt.Println("Warning: Could not determine compiled_rules.txt cache path from AdGuardHome.yaml; skipping
767 cache update.")
768 }
769
770 adguardReachable := false
771
772 dockerComposeYamlExists := false
773 if _, err := os.Stat("docker-compose.yml"); err == nil {
774     dockerComposeYamlExists = true
775 }
776
777 dockerPath, err := exec.LookPath("docker")
778 if err == nil && dockerComposeYamlExists {
779     cmd := exec.Command(dockerPath, "compose", "ps", "--status", "running", "-q", "adguardhome")
780     out, err := cmd.Output()
781     if err == nil && strings.TrimSpace(string(out)) != "" {
782         fmt.Println("Force-recreating AdGuard Home container to drop persisted filter state...")
783         upCmd := exec.Command(dockerPath, "compose", "up", "-d", "--force-recreate", "--no-deps", "adguardhome")
784         err := upCmd.Run()
785         if err == nil {
786             fmt.Println("AdGuard Home force-recreated successfully.")
787             time.Sleep(8 * time.Second)
788             adguardReachable = true
789         } else {
790             fmt.Println("Failed to restart AdGuard Home. Is it running?")
791         }
792     }
793 }
794 }
795 }

```

```

791 if adguardReachable {
792     agPass := "nullexit"
793     if bytesData, err := os.ReadFile(".env"); err == nil {
794         scanner := bufio.NewScanner(strings.NewReader(string(bytesData)))
795         for scanner.Scan() {
796             line := scanner.Text()
797             if strings.HasPrefix(line, "ADGUARD_PASSWORD=") {
798                 parts := strings.SplitN(line, "=", 2)
799                 if len(parts) == 2 {
800                     agPass = strings.Trim(strings.TrimSpace(parts[1]), "\\\"")
801                 }
802                 break
803             }
804         }
805     }
806
807     creds := base64.StdEncoding.EncodeToString([]byte(fmt.Sprintf("admin:%s", agPass)))
808     req, err := http.NewRequest("POST", "http://127.0.0.1:3000/control/filtering/refresh", bytes.NewBuffer([]
809         byte("{}")))
810     if err == nil {
811         req.Header.Set("Authorization", fmt.Sprintf("Basic %s", creds))
812         req.Header.Set("Content-Type", "application/json")
813         client := &http.Client{Timeout: 15 * time.Second}
814         resp, err := client.Do(req)
815         if err == nil {
816             fmt.Printf("Triggered AdGuard filter refresh via REST API (HTTP=%d).\n", resp.StatusCode)
817             resp.Body.Close()
818         } else {
819             fmt.Printf("Warning: AdGuard filter refresh API call failed (%v). New whitelist will not load until
820                 next refresh tick (~24h) or manual refresh.\n", err)
821         }
822     }
823 }
824
825 compileIpBlocklist()
826 }

```

## 44 scripts/setup-common.sh

```

1  #!/bin/bash
2  # 1. Docker
3  step "Checking Docker"
4
5  if ! command -v docker >> output.log 2>&1; then
6      echo ""
7      if [[ "$OSTYPE" == "darwin"* ]]; then
8          echo " Docker is not installed."
9          # Note: Homebrew aggressively rolls forward and discourages version pinning.
10         # This setup is confirmed stable on Colima v0.10.3 and Docker v29.6.1.
11         read -rp " Would you like to automatically install Colima & Docker via Homebrew? [y/N]: "
12         INSTALL_DOCKER
13         if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
14             step "Installing Colima and Docker..."
15             brew install colima docker docker-compose
16             step "Starting Colima VM..."
17             colima start --memory 0.6 --vm-type vz --network-address --network-mode bridged
18         else
19             echo ""
20             echo " Please install Docker manually. Two options:"
21             echo " A) Colima (Recommended): brew install colima docker docker-compose && colima start --memory
22                 0.6 --vm-type vz --network-address --network-mode bridged"
23             echo " B) Docker Desktop: https://www.docker.com/products/docker-desktop/"
24             die "Install Docker, then re-run this script."
25         fi
26     else
27         echo " Docker is not installed."
28         # Note: The Docker convenience script always fetches the latest stable release.
29         # This setup is confirmed stable on Docker Engine v29.6.1.
30         read -rp " Would you like to automatically install Docker? (Requires sudo) [y/N]: " INSTALL_DOCKER
31         if [[ "$INSTALL_DOCKER" == "y" || "$INSTALL_DOCKER" == "Y" ]]; then
32             step "Installing Docker..."
33             curl -fsSL https://get.docker.com | sh
34             sudo usermod -aG docker "$USER"
35             echo -e "\n\033[0;32m[OK] Docker installed successfully!\033[0m"
36             echo -e "\033[0;33m[!] CRITICAL: You must log out of your SSH session and log back in for Docker
37                 permissions to apply.\033[0m"

```

```

35         die "Please log out, log back in, and re-run setup.sh."
36     else
37         echo "   Please install Docker manually:"
38         echo "       curl -fsSL https://get.docker.com | sh"
39         echo "       sudo usermod -aG docker \$USER && newgrp docker"
40         die "Install Docker, then re-run this script."
41     fi
42 fi
43 fi
44
45 if ! docker info >> output.log 2>&1; then
46     echo ""
47     if [[ "$OSTYPE" == "darwin"* ]]; then
48         echo "   Docker is installed but not running."
49         echo "   If using Colima:          colima start --memory 0.6 --vm-type vz --network-address --network-mode
bridged"
50         echo "   If using Docker Desktop: open the app."
51     else
52         echo "   Docker is installed but not running."
53         echo "   Run: sudo systemctl start docker"
54     fi
55     die "Start Docker, then re-run this script."
56 fi
57
58 if ! docker compose version >> output.log 2>&1; then
59     die "Docker Compose v2 not found. Update Docker or install the compose plugin:\n  https://docs.docker.com/
compose/install/"
60 fi
61
62 ok "Docker \$(docker --version | grep -oE '[0-9.]+' | head -1) is running."
63
64 # 1.5 Python 3
65 step "Checking Python 3"
66
67 if ! command -v python3 >> output.log 2>&1; then
68     echo ""
69     if [[ "$OSTYPE" == "darwin"* ]]; then
70         echo "   Python 3 is not installed."
71         echo "   macOS includes it via Xcode Command Line Tools, or you can install it via Homebrew:"
72         echo "       brew install python3"
73         die "Install Python 3, then re-run this script."
74     else
75         echo "   Python 3 is not installed."
76         echo "   Please install it using your system's package manager. For example:"
77         echo "       Debian/Ubuntu:  sudo apt install python3"
78         echo "       Fedora:         sudo dnf install python3"
79         echo "       Arch:           sudo pacman -S python"
80         die "Install Python 3, then re-run this script."
81     fi
82 fi
83
84 # Note: This setup is confirmed stable on Python 3.9.6 (macOS default).
85 PYTHON_VER=\$(python3 --version 2>&1 | head -n 1)
86 ok "\$PYTHON_VER is installed."
87
88 # 2. Tailscale (host)
89 step "Checking Tailscale (host)"
90 # The host needs its own Tailscale installation separate from the Docker
91 # container so Toggle-Gateway can run 'tailscale down' before stopping
92 # containers and 'tailscale up' after starting them. Without this, the
93 # exit-node deadlock that setup.sh is designed to prevent can still occur.
94
95 # Resolve the tailscale binary: CLI (brew/package) or bundled macOS app
96 TAILSCALE_BIN=""
97 if command -v tailscale >> output.log 2>&1; then
98     TAILSCALE_BIN="tailscale"
99 elif [[ -x "/Applications/Tailscale.app/Contents/MacOS/Tailscale" ]]; then
100     TAILSCALE_BIN="/Applications/Tailscale.app/Contents/MacOS/Tailscale"
101 fi
102
103 RECOMMENDED_TS_VERSION="1.98.5"
104
105 if [[ -z "\$TAILSCALE_BIN" ]]; then
106     warn "Tailscale not found on host installing..."
107
108     if [[ "$OSTYPE" == "darwin"* ]]; then
109         if command -v brew >> output.log 2>&1; then
110             step "Installing Tailscale CLI formula via Homebrew (standalone, no .app GUI)..."
111             # Note: We install Tailscale via Homebrew, targeting the stable version (recommended: 1.98.5)
112             brew install tailscale -q
113             TAILSCALE_BIN="tailscale"

```

```

114         step "Starting tailscaled as a per-user LaunchAgent (auto-starts at login)..."
115         if brew services start tailscale 2>&1 | grep -qE 'Successfully|started'; then
116             ok "tailscaled LaunchAgent registered."
117         else
118             warn "Could not auto-start tailscaled. Run 'brew services start tailscale' manually."
119             warn "For a system-wide LaunchDaemon that runs before login, run instead:"
120             warn "    sudo brew services start tailscale"
121         fi
122         echo ""
123     else
124         die "Homebrew not found. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION}):
125         }\n https://tailscale.com/download/mac\n Then re-run this script."
126     fi
127
128     elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
129         curl -fsSL https://tailscale.com/install.sh | sh
130         sudo systemctl enable --now tailscaled 2>> output.log || true
131         TAILSCALE_BIN="tailscale"
132
133     else
134         die "Unsupported OS. Install Tailscale manually (recommended version: ${RECOMMENDED_TS_VERSION}):
135         }\n https://tailscale.com/download\n Then re-run this script."
136     fi
137
138     TAILSCALE_VER=$((${TAILSCALE_BIN} version 2>> output.log | head -n 1)
139     if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
140         warn "Tailscale installed, but version ($TAILSCALE_VER) differs from recommended version (
141         $RECOMMENDED_TS_VERSION)."
142     else
143         ok "Tailscale installed (version $TAILSCALE_VER)."
144     fi
145
146     TAILSCALE_VER=$((${TAILSCALE_BIN} version 2>> output.log | head -n 1)
147     if [[ "$TAILSCALE_VER" != "$RECOMMENDED_TS_VERSION" ]]; then
148         warn "Tailscale is already installed ($TAILSCALE_VER), but recommended version is
149         $RECOMMENDED_TS_VERSION for host/container consistency."
150     else
151         ok "Tailscale is already installed at recommended version ($TAILSCALE_VER)."
152     fi
153 fi
154
155 # Authenticate the host machine if it isn't already connected.
156 # This is an interactive step it opens a browser login URL.
157 # We do it now, before containers start, so Toggle-Gateway can
158 # safely run 'tailscale down/up' from day one.
159 if ! ${TAILSCALE_BIN} status >> output.log 2>&1; then
160     echo ""
161     echo " Your host machine needs to join your Tailscale network."
162     echo " A browser window or login URL will appear authenticate, then"
163     echo " return here. The script will continue automatically."
164     echo ""
165     if [[ "$OSTYPE" == "linux-gnu"* ]]; then
166         sudo ${TAILSCALE_BIN} up
167     else
168         ${TAILSCALE_BIN} up
169     fi
170     ok "Host machine connected to Tailscale."
171 else
172     ok "Host machine is already connected to Tailscale."
173 fi
174
175 # 3. wgcf
176 step "Checking wgcf (Cloudflare WARP key tool)"
177
178 if ! command -v wgcf >> output.log 2>&1; then
179     warn "wgcf not found installing..."
180
181     if [[ "$OSTYPE" == "darwin"* ]]; then
182         command -v brew >> output.log 2>&1 || die "Homebrew not found.\nInstall wgcf manually: https://github.
183         com/ViRb3/wgcf/releases"
184         brew install wgcf -q
185     elif [[ "$OSTYPE" == "linux-gnu"* ]]; then
186         ARCH=$(uname -m)
187         case $ARCH in
188             x86_64) WGCF_ARCH="amd64" ;;
189             aarch64) WGCF_ARCH="arm64" ;;
190             armv7l) WGCF_ARCH="armv7" ;;
191             *) die "Unsupported architecture: $ARCH.\nInstall wgcf manually: https://github.com/ViRb3/wgcf/
192             releases" ;;

```



```

189     esac
190
191     echo " Fetching latest release..."
192     WGCF_VERSION=$(curl -sf https://api.github.com/repos/ViRb3/wgcf/releases/latest \
193     | grep 'tag_name' | cut -d'"' -f4)
194     [[ -z "$WGCF_VERSION" ]] && die "Could not fetch wgcf version. Check your internet connection."
195
196     sudo curl -sfl \
197     "https://github.com/ViRb3/wgcf/releases/download/${WGCF_VERSION}/wgcf_${WGCF_VERSION#v}_linux_${
198     WGCF_ARCH}" \
199     -o /usr/local/bin/wgcf
200     sudo chmod +x /usr/local/bin/wgcf
201 else
202     die "Unsupported OS. Install wgcf manually: https://github.com/ViRb3/wgcf/releases"
203 fi
204
205 ok "wgcf is ready."
206
207 # 4. WARP profile
208 step "Generating Cloudflare WARP profile"
209
210 if [[ -f "wgcf-profile.conf" ]]; then
211     warn "wgcf-profile.conf already exists skipping key generation."
212 else
213     wgcf register --accept-tos 2>> output.log || die "wgcf register failed. Check your internet connection."
214     wgcf generate 2>> output.log || die "wgcf generate failed."
215     ok "WARP profile generated."
216 fi
217
218 # Parse keys (IPv4 address only gluetun doesn't need the IPv6 one)
219 PRIVATE_KEY=$(awk '/PrivateKey/{print $3}' wgcf-profile.conf)
220 PUBLIC_KEY=$(awk '/PublicKey/{print $3}' wgcf-profile.conf)
221 ADDRESSES=$(awk '/Address/{print $3}' wgcf-profile.conf | grep -v ':' | head -1)
222
223 [[ -z "$PRIVATE_KEY" ]] && die "Could not parse PrivateKey from wgcf-profile.conf"
224 [[ -z "$PUBLIC_KEY" ]] && die "Could not parse PublicKey from wgcf-profile.conf"
225 [[ -z "$ADDRESSES" ]] && die "Could not parse IPv4 Address from wgcf-profile.conf"
226
227 ok "WARP keys extracted."
228
229 # 5. Tailscale auth key
230 step "Tailscale authentication"
231
232 TS_AUTHKEY=""
233 if [[ -f ".env" ]]; then
234     EXISTING_TS_KEY=$(read_env_var TS_AUTHKEY)
235     if [[ -n "$EXISTING_TS_KEY" ]]; then
236         TS_AUTHKEY="$EXISTING_TS_KEY"
237         warn "Found existing TS_AUTHKEY in .env. Skipping prompt."
238     fi
239 fi
240
241 if [[ -z "$TS_AUTHKEY" ]]; then
242     echo ""
243     echo " Generate a key at: https://login.tailscale.com/admin/settings/keys"
244     echo " 'Generate auth key' enable Reusable if you plan to re-run setup"
245     echo ""
246     read -rp " Paste your Tailscale auth key: " TS_AUTHKEY
247     echo ""
248 fi
249
250 [[ -z "$TS_AUTHKEY" ]] && die "Tailscale auth key is required."
251 [[ "$TS_AUTHKEY" != tskey-auth-* ]] && warn "Key doesn't look like a Tailscale auth key proceeding anyway."
252 ok "Auth key accepted."
253
254 # 6. AdGuard Home admin password
255 step "AdGuard Home admin account"
256
257 # Hardcoded default credentials to prevent lockout
258 # Username: admin
259 # Password: nullexit
260 if [[ -f ".env" ]]; then
261     EXISTING_AG_PASS=$(read_env_var ADGUARD_PASSWORD)
262     if [[ -n "$EXISTING_AG_PASS" ]]; then
263         ADGUARD_PASSWORD="$EXISTING_AG_PASS"
264         ADGUARD_PWD_ESC="$EXISTING_AG_PASS"
265         ok "Found existing ADGUARD_PASSWORD in .env."
266     else
267         ADGUARD_PASSWORD="nullexit"
268         ADGUARD_PWD_ESC="nullexit"

```

```

269         ok "Default password set to 'nullexit'."
270     fi
271 else
272     ADGUARD_PASSWORD="nullexit"
273     ADGUARD_PWD_ESC="nullexit"
274     ok "Default password set to 'nullexit'."
275 fi
276
277 # 7. Write .env
278 step "Writing .env"
279
280 if [[ -f ".env" ]]; then
281     read -rp " .env already exists. Overwrite? [y/N]: " OVERWRITE
282     if [[ "$OVERWRITE" != "y" && "$OVERWRITE" != "Y" ]]; then
283         warn "Keeping existing .env."
284     else
285         WRITE_ENV=1
286     fi
287 else
288     WRITE_ENV=1
289 fi
290
291 if [[ "${WRITE_ENV:-0}" == "1" ]]; then
292     # Preserve existing configuration flags if .env already exists
293     EXISTING_PROFILE="medium"
294     EXISTING_MSS="1120"
295     EXISTING_HIJACK="true"
296     EXISTING_EXIT="true"
297     EXISTING_THRESH="6"
298     EXISTING_KILL="false"
299     EXISTING_BLOCKED="kp il"
300     EXISTING_HOST_MTU="1200"
301
302     if [[ -f ".env" ]]; then
303         [[ -n "$(read_env_var GATEWAY_RULE_PROFILE)" ]] && EXISTING_PROFILE=$(read_env_var GATEWAY_RULE_PROFILE)
304     )
305     [[ -n "$(read_env_var GATEWAY_MSS)" ]] && EXISTING_MSS=$(read_env_var GATEWAY_MSS)
306     [[ -n "$(read_env_var GATEWAY_HIJACK_HOST)" ]] && EXISTING_HIJACK=$(read_env_var GATEWAY_HIJACK_HOST)
307     [[ -n "$(read_env_var GATEWAY_USE_EXIT_NODE)" ]] && EXISTING_EXIT=$(read_env_var GATEWAY_USE_EXIT_NODE)
308     [[ -n "$(read_env_var WARP_FAIL_THRESHOLD)" ]] && EXISTING_THRESH=$(read_env_var WARP_FAIL_THRESHOLD)
309     [[ -n "$(read_env_var KILL_SWITCH)" ]] && EXISTING_KILL=$(read_env_var KILL_SWITCH)
310     [[ -n "$(read_env_var BLOCKED_COUNTRIES)" ]] && EXISTING_BLOCKED=$(read_env_var BLOCKED_COUNTRIES)
311     [[ -n "$(read_env_var HOST_MTU)" ]] && EXISTING_HOST_MTU=$(read_env_var HOST_MTU)
312 fi
313
314 cat > .env <<EOF
315 WIREGUARD_PRIVATE_KEY=${PRIVATE_KEY}
316 WIREGUARD_PUBLIC_KEY=${PUBLIC_KEY}
317 WIREGUARD_ADDRESSES=${ADDRESSES}
318 TS_AUTHKEY=${TS_AUTHKEY}
319 GATEWAY_RULE_PROFILE=${EXISTING_PROFILE}
320 # Safe MSS for double-tunneled traffic (Tailscale + WARP). Change to 1180 if you experience slow speeds and
321 # have a healthy path.
322 GATEWAY_MSS=${EXISTING_MSS}
323 # Lower the host Tailscale MTU to 1200 to prevent fragmentation when passing through the 1280-byte WARP tunnel.
324 HOST_MTU=${EXISTING_HOST_MTU}
325 # Set to false to run as a 'Headless Server' (Docker acts as an exit node, but the host Mac/Linux's own
326 # internet is NOT hijacked).
327 GATEWAY_HIJACK_HOST=${EXISTING_HIJACK}
328 GATEWAY_USE_EXIT_NODE=${EXISTING_EXIT}
329 WARP_FAIL_THRESHOLD=${EXISTING_THRESH}
330 KILL_SWITCH=${EXISTING_KILL}
331 ADGUARD_PASSWORD=${ADGUARD_PASSWORD}
332 BLOCKED_COUNTRIES="${EXISTING_BLOCKED}"
333 EOF
334
335 ok ".env written."
336 fi
337
338 # 8. Prepare directories
339 step "Preparing AdGuard directories"
340
341 mkdir -p adguard/conf adguard/work/userfilters
342
343 # Remove empty AdGuardHome.yaml that would cause a crash loop (same logic as Toggle-Gateway scripts)
344 if [[ -f "adguard/conf/AdGuardHome.yaml" && ! -s "adguard/conf/AdGuardHome.yaml" ]]; then
345     rm "adguard/conf/AdGuardHome.yaml"
346     warn "Removed empty AdGuardHome.yaml to prevent crash loop."
347 fi
348
349 ok "Directories ready."

```

```

347
348 # 9. Compile DNS filter rules
349 step "Compiling DNS filter rules (this may take a minute)"
350
351
352 # 10. Start containers
353 step "Starting containers"
354 # Note: tailscale depends on adguardhome being healthy (port 5335 up),
355 # so Docker Compose will hold it back automatically until the wizard is done.
356 docker compose up -d
357 ok "Containers starting."
358
359 # 11. Wait for AdGuard Home setup wizard
360 step "Waiting for AdGuard Home setup wizard"
361 echo " (up to 60 seconds...)"
362
363 MAX=60; ELAPSED=0
364 until docker exec curl -sf http://127.0.0.1:3000 >> output.log 2>&1; do
365     sleep 2; ELAPSED=$((ELAPSED + 2))
366     [[ $ELAPSED -ge $MAX ]] && die "AdGuard Home didn't come up.\nDebug with: docker compose logs adguardhome"
367 done
368
369 ok "AdGuard Home is up."
370
371 # 12. Configure AdGuard Home via API
372 step "Configuring AdGuard Home"
373
374 # Run all API calls in a single docker exec sh -c so the session cookie file
375 # is shared across calls without leaving the container.
376 docker exec routing-fix sh -c "
377     set -e
378
379     # Step 1: Complete the setup wizard (sets admin creds + DNS port to 5335)
380     curl -sf -X POST http://127.0.0.1:3000/control/install/configure \
381         -H 'Content-Type: application/json' \
382         -d '{"web":{"ip":"0.0.0.0","port":3000},"dns":{"ip":"0.0.0.0","port":5335},"username":"admin","password":"'${ADGUARD_PWD_ESC}'" \
383         >/dev/null || true
384
385     # Step 2: Wait for AdGuard to restart after wizard
386     sleep 4
387     WAITED=0
388     until curl -sf http://127.0.0.1:3000 >/dev/null 2>&1; do
389         sleep 2; WAITED=$((WAITED + 2))
390         [ $WAITED -ge 30 ] && { echo 'AdGuard did not restart in time'; exit 1; }
391     done
392
393     # Step 3: Login to get session cookie
394     curl -sf -X POST http://127.0.0.1:3000/control/login \
395         -H 'Content-Type: application/json' \
396         -d '{"name":"admin","password":"'${ADGUARD_PWD_ESC}'" \
397         -c /tmp/agh.cookie >/dev/null
398
399     # Step 4: Point upstream DNS back through the WARP tunnel
400     curl -sf -X POST http://127.0.0.1:3000/control/dns_config \
401         -H 'Content-Type: application/json' \
402         -b /tmp/agh.cookie \
403         -d '{"upstream_dns":["127.0.0.1:53","[/onion/]127.0.0.1:5353"],"bootstrap_dns":["1.1.1.1","8.8.8.8"],"upstream_mode":"load_balance"}' \
404         >/dev/null
405
406     # Step 5: Register the compiled rules file as a filter list
407     curl -sf -X POST http://127.0.0.1:3000/control/filtering/add_url \
408         -H 'Content-Type: application/json' \
409         -b /tmp/agh.cookie \
410         -d '{"name":"Custom Compiled Rules","url":"/opt/adguardhome/work/userfilters/compiled_rules.txt"}' \
411         >/dev/null || true
412
413     # Step 6: Force a filter refresh to load the rules immediately
414     curl -sf -X POST http://127.0.0.1:3000/control/filtering/refresh \
415         -H 'Content-Type: application/json' \
416         -b /tmp/agh.cookie \
417         -d '{"whitelist":false}' \
418         >/dev/null || true
419
420     rm -f /tmp/agh.cookie
421 "
422
423 ok "AdGuard Home configured."
424

```

```

425 # 13. Wait for Tailscale
426 step "Waiting for Tailscale to authenticate"
427 echo " (Tailscale only starts once AdGuard's healthcheck passes up to 90s)"
428
429 MAX=90; ELAPSED=0; TS_IP=""
430 until [[ -n "$TS_IP" ]]; do
431     TS_IP=$(docker exec tailscale tailscale ip -4 2>> output.log || true)
432     if [[ -z "$TS_IP" ]]; then
433         sleep 3; ELAPSED=$((ELAPSED + 3))
434         if [[ $ELAPSED -ge $MAX ]]; then
435             warn "Tailscale hasn't authenticated yet."
436             warn "Once it does, re-run the toggle script it will resolve the IP dynamically."
437             break
438         fi
439     fi
440 done
441
442 if [[ -n "$TS_IP" ]]; then
443     ok "Tailscale IP: ${TS_IP} (resolved dynamically by the toggle script on each startup)"
444 fi

```

## 45 scripts/sweep.sh

```

1 #!/bin/bash
2 # scripts/sweep.sh nullexit gateway health sweep
3 #
4 # Automates the 6-check post-restart health sweep (sweep.md 1 + tailnet ping).
5 # It answers one question: "after a toggle/restart/wake, is the gateway
6 # actually healthy end-to-end no leak, direct P2P, DNS filtering live,
7 # kill-switch armed?" It is READ-ONLY: it never mutates routing, PF, DNS,
8 # or containers, so it is safe to run against a live gateway at any time
9 # (including from a remote session see the caveat in agent.md L16).
10 #
11 # The 6 checks (1-5 mirror sweep.md 1):
12 # 1. Containers warp, adguardhome, tailscale, tor, routing-fix up/healthy
13 # 2. Double-tunnel/leak container WARP exit IP == host exit IP, both warp=on
14 # 3. Hostcontainer path Tailscale exit-node is a DIRECT P2P conn, not DERP relay
15 # 4. AdGuard filtering compiled rule counts present + live DNS interception
16 # 5. PF kill-switch pfctl Enabled + anchor com.apple/nullexit carries rules
17 # 6. Tailnet reachability every ONLINE Tailscale peer answers a tailscale ping
18 #
19 # A timestamped report is written to the project root (one file per run, so
20 # runs can be diffed). stdout carries a coloured summary + a final PASS/FAIL
21 # verdict; exit code is 0 when every check passes, 1 otherwise (so it can gate
22 # CI / launchd / a post-restart hook).
23 #
24 # Usage:
25 # bash scripts/sweep.sh # run the sweep, write report, print verdict
26 # bash scripts/sweep.sh --quiet # only the final verdict line + non-zero on FAIL
27 # bash scripts/sweep.sh --help # this message
28
29 set -uo pipefail
30 # NOTE: errexit is intentionally OFF. Several probes are EXPECTED to fail on an
31 # unhealthy gateway and that failure IS the diagnostic signal we record each
32 # exit status rather than letting one abort the whole sweep.
33
34 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
35 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
36 TIMESTAMP="$(date -u +%Y%m%dT%H%M%SZ)"
37 cd "$PROJECT_ROOT" || { echo "[FATAL] cannot cd to $PROJECT_ROOT" >&2; exit 1; }
38 source "$SCRIPT_DIR/common.sh"
39 LOG_FILE="$PROJECT_ROOT/output.log"
40 REPORT_FILE="$PROJECT_ROOT/sweep-$TIMESTAMP.txt"
41
42 # Argument parsing
43 QUIET=false
44 case "${1:-}" in
45     --quiet|-q) QUIET=true ;;
46     --help|-h) sed -n '2,27p' "$0"; exit 0 ;;
47     *)
48         : ;;
49     *)
50         echo "Unknown argument: $1" >&2
51         echo "Run with --help for usage." >&2
52         exit 2 ;;
53 esac
54
55 # Colours (auto-disabled when stdout is not a tty, e.g. piped)

```

```

54 if [ -t 1 ]; then
55     RED='\033[0;31m'; GREEN='\033[0;32m'; YELLOW='\033[1;33m'
56     BOLD='\033[1m'; NC='\033[0m'
57 else
58     RED=''; GREEN=''; YELLOW=''; BOLD=''; NC=''
59 fi
60
61 # Homebrew on PATH (same as toggle.sh / recover.sh / diagnose-host-leak)
62 export PATH="/opt/homebrew/bin:/usr/local/bin:$PATH"
63
64 # Output helpers write to BOTH stdout and the report file
65 exec 3>> "$REPORT_FILE" || {
66     printf '\n[FATAL] cannot open %s for write\n' "$REPORT_FILE" >&2
67     exit 1
68 }
69 emit() { [ "$QUIET" = true ] || printf '%b\n' "$1"; printf '%b\n' "$1" | sed 's/\x1b\[[0-9;]*m//g' >&3; }
70 section() { emit ""; emit "${BOLD} $1 ${NC}"; }
71 line() { emit " $1"; }
72 ok() { line "${GREEN}${NC} $1"; }
73 warn() { line "${YELLOW}${NC} $1"; }
74 fail() { line "${RED}${NC} $1"; }
75
76 # common.sh redefines run_with_timeout; it is a pure-bash macOS-safe timeout
77 # (GNU `timeout` is absent on stock macOS). We reuse the sourced version.
78
79 # Result accounting
80 # Each check appends "PASS"/"WARN"/"FAIL" so the verdict is a pure roll-up.
81 declare -a RESULTS=(
82 record() { RESULTS+=("${1}|${2}"); } # status|label
83
84 # Header
85 if [ "$QUIET" != true ]; then
86     emit ""
87     emit ""
88     emit " n u l l e x i t   h e a l t h   s w e e p "
89     emit ""
90     emit " Timestamp (UTC): $TIMESTAMP"
91     emit " Project root:     $PROJECT_ROOT"
92     emit " Report file:      $REPORT_FILE"
93     emit ""
94 fi
95
96 GATEWAY_TS_IP="$(read_adguard_ip)" # gateway Tailscale IP from .gateway_ip
97
98 #
99 section "1/6 Containers"
100 #
101 # All five services must be running; warp + adguardhome additionally expose a
102 # healthcheck and must be 'healthy'. `docker compose ps` is the source of truth.
103 EXPECTED_SERVICES="warp adguardhome tailscale tor routing-fix"
104 if ! command -v docker >/dev/null 2>&1; then
105     fail "docker CLI not on PATH cannot inspect containers"
106     record FAIL "containers"
107 else
108     PS_OUT=$(run_with_timeout 20 docker compose ps 2>>"$LOG_FILE")
109     missing=""; unhealthy=""
110     for svc in $EXPECTED_SERVICES; do
111         row=$(printf '%s\n' "$PS_OUT" | awk -v s="$svc" '1==s || $0 ~ ("[:space:]"s"[:space:]")')
112         if [ -z "$row" ] || ! printf '%s' "$row" | grep -qiE 'up|running'; then
113             missing="$missing $svc"
114         elif printf '%s' "$row" | grep -qi 'unhealthy'; then
115             unhealthy="$unhealthy $svc"
116         fi
117     done
118     if [ -z "$missing" ] && [ -z "$unhealthy" ]; then
119         ok "all 5 services up (warp/adguardhome healthy, tailscale/tor/routing-fix running)"
120         record PASS "containers"
121     else
122         [ -n "$missing" ] && fail "not up/running:$missing"
123         [ -n "$unhealthy" ] && fail "reporting unhealthy:$unhealthy"
124         record FAIL "containers"
125     fi
126     line " docker compose ps "
127     printf '%s\n' "$PS_OUT" | awk '{print "    "$0}' >&3
128     [ "$QUIET" = true ] || printf '%s\n' "$PS_OUT" | awk '{print "    "$0}'
129 fi
130
131 #
132 section "2/6 Double-tunnel / IP-leak"
133 #
134 # The container's WARP egress and the host's egress must be the SAME public IP

```

```

135 # with warp=on for both. Equal IPs all host traffic is exiting through the
136 # same WARP tunnel the container uses (no split / raw-ISP leak).
137 # The warp container (gluetun) ships wget, not curl fetch the trace from
138 # inside its netns with whichever client is present.
139 CTR_TRACE=$(run_with_timeout 12 docker compose exec -T warp sh -c \
140     'curl -s --max-time 8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null \
141     || wget -qO- --timeout=8 https://www.cloudflare.com/cdn-cgi/trace 2>/dev/null' \
142     2>>"$LOG_FILE")
143 HOST_TRACE=$(run_with_timeout 12 curl -s --max-time 8 \
144     https://www.cloudflare.com/cdn-cgi/trace 2>>"$LOG_FILE")
145 CTR_IP=$(printf '%s' "$CTR_TRACE" | awk -F= '/^ip={print $2;exit}')
146 CTR_WARP=$(printf '%s' "$CTR_TRACE" | awk -F= '/^warp={print $2;exit}')
147 HOST_IP=$(printf '%s' "$HOST_TRACE" | awk -F= '/^ip={print $2;exit}')
148 HOST_WARP=$(printf '%s' "$HOST_TRACE" | awk -F= '/^warp={print $2;exit}')
149 line "container exit: ip=${CTR_IP:-?} warp=${CTR_WARP:-?}"
150 line "host exit: ip=${HOST_IP:-?} warp=${HOST_WARP:-?}"
151 if [ -z "$CTR_IP" ] || [ -z "$HOST_IP" ]; then
152     fail "could not obtain both egress IPs (container or host trace failed)"
153     record FAIL "leak"
154 elif [ "$CTR_IP" = "$HOST_IP" ] && [ "$HOST_WARP" = "on" ]; then
155     ok "no leak host egress == container egress ($HOST_IP), warp=on"
156     record PASS "leak"
157 elif [ "$HOST_WARP" != "on" ]; then
158     fail "host warp=$HOST_WARP host traffic is NOT going through WARP (LEAK)"
159     record FAIL "leak"
160 else
161     fail "egress mismatch: host $HOST_IP != container $CTR_IP split/leak"
162     record FAIL "leak"
163 fi
164
165 #
166 section "3/6 Host container path (direct P2P, not DERP relay)"
167 #
168 # `tailscale status` annotates each peer with 'direct <ip:port>' or
169 # 'relay <derp>'. The gateway peer should be DIRECT a DERP relay hop adds
170 # latency and signals NAT traversal failure between host and gateway.
171 if ! command -v tailscale >/dev/null 2>&1; then
172     fail "tailscale CLI not on PATH"
173     record FAIL "p2p"
174 else
175     TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
176     PEER_LINE=$(printf '%s\n' "$TS_OUT" | grep -E "(^|[:space:])$GATEWAY_TS_IP([[:space:]]|$)" | head -1)
177     if [ -z "$PEER_LINE" ]; then
178         warn "gateway $GATEWAY_TS_IP not found in host peer list"
179         line "$(printf '%s\n' "$TS_OUT" | head -5 | awk '{print "    "$0}')"
180         record WARN "p2p"
181     elif printf '%s' "$PEER_LINE" | grep -q 'direct '; then
182         conn=$(printf '%s' "$PEER_LINE" | grep -oE 'direct [0-9.]+:[0-9.]+')
183         ok "direct P2P to gateway $conn (no DERP relay)"
184         record PASS "p2p"
185     elif printf '%s' "$PEER_LINE" | grep -q 'relay '; then
186         relay=$(printf '%s' "$PEER_LINE" | grep -oE 'relay "[^"]+"')
187         warn "gateway reached via DERP $relay (not direct) NAT traversal degraded"
188         record WARN "p2p"
189     else
190         warn "gateway peer state indeterminate (idle/no active conn)"
191         line "    $PEER_LINE"
192         record WARN "p2p"
193     fi
194 fi
195
196 #
197 section "4/6 AdGuard filtering"
198 #
199 # Two signals: (a) the compiled ruleset exists with a non-trivial block count,
200 # (b) the gateway resolver actually answers a query (live interception).
201 COMPILED="adguard/work/userfilters/compiled_rules.txt"
202 if [ -f "$COMPILED" ]; then
203     TOTAL=$(grep -m1 '! Total Block Rules:' "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
204     NATIVE=$(grep -m1 '! Native AdGuard Rules:' "$COMPILED" | awk -F': ' '{print $2}' | tr -d '[:space:]')
205     if [ -n "$TOTAL" ] && [ "$TOTAL" -gt 0 ] 2>/dev/null; then
206         ok "compiled ruleset present Total Block Rules: $TOTAL / Native AdGuard: ${NATIVE:-?}"
207         RULES_OK=true
208     else
209         warn "compiled_rules.txt present but block-rule count unreadable"
210         RULES_OK=false
211     fi
212 else
213     warn "compiled_rules.txt missing (rule-compile may not have run)"
214     RULES_OK=false
215 fi

```

```

216 # Live interception: the gateway resolver must answer (dig via .gateway_ip).
217 DNS_OK=false
218 if command -v dig >/dev/null 2>&1 && [ -n "$GATEWAY_TS_IP" ]; then
219     DIG_OUT=$(run_with_timeout 8 dig +time=3 +tries=1 @$GATEWAY_TS_IP google.com A 2>>"$LOG_FILE")
220     if printf '%s' "$DIG_OUT" | grep -qE 'ANSWER SECTION|status: NOERROR'; then
221         ok "live DNS interception gateway $GATEWAY_TS_IP resolved google.com"
222         DNS_OK=true
223     else
224         warn "gateway $GATEWAY_TS_IP did not answer a test query"
225     fi
226 else
227     warn "skipped live-resolution probe (dig missing or gateway IP unknown)"
228 fi
229 if [ "$RULES_OK" = true ] && [ "$DNS_OK" = true ]; then
230     record PASS "adguard"
231 elif [ "$RULES_OK" = true ] || [ "$DNS_OK" = true ]; then
232     record WARN "adguard"
233 else
234     record FAIL "adguard"
235 fi
236
237 #
238 section "5/6 PF kill-switch"
239 #
240 # PF must be Enabled AND the com.apple/nullexit anchor must carry rules.
241 # An empty/absent anchor means the fail-closed lane is not actually armed.
242 PF_STATUS=$(run_with_timeout 6 sudo -n pfctl -s info 2>>"$LOG_FILE" | awk '/^Status:/{print $2;exit}')
243 ANCHOR_RULES=$(run_with_timeout 6 sudo -n pfctl -a com.apple/nullexit -sr 2>>"$LOG_FILE")
244 if [ "$PF_STATUS" = "Enabled" ] && [ -n "$ANCHOR_RULES" ]; then
245     RULE_N=$(printf '%s\n' "$ANCHOR_RULES" | grep -cE '\s*')
246     ok "PF Enabled + anchor com.apple/nullexit armed ($RULE_N rules)"
247     line "$(printf '%s\n' "$ANCHOR_RULES" | head -4 | awk '{print "      "$0}')"
248     record PASS "pf"
249 elif [ "$PF_STATUS" = "Enabled" ]; then
250     fail "PF Enabled but anchor com.apple/nullexit has NO rules kill-switch not armed"
251     record FAIL "pf"
252 elif [ -z "$PF_STATUS" ]; then
253     warn "could not read PF status (needs passwordless sudo for 'pfctl -s info')"
254     record WARN "pf"
255 else
256     fail "PF Status: $PF_STATUS kill-switch disabled"
257     record FAIL "pf"
258 fi
259
260 #
261 section "6/6 Tailnet reachability (ping all online devices)"
262 #
263 # Every peer the coordination server reports as ONLINE should answer a
264 # `tailscale ping` (TSMP probes the WireGuard data path; ICMP is avoided since
265 # phones/laptops often drop it). Offline peers and this host are skipped.
266 # An online-but-unreachable peer means control plane and data path disagree.
267 if ! command -v tailscale >/dev/null 2>&1; then
268     fail "tailscale CLI not on PATH"
269     record FAIL "tailnet"
270 else
271     SELF_TS_IP=$(run_with_timeout 6 tailscale ip -4 2>>"$LOG_FILE")
272     [ -n "$TS_OUT:-" ] || TS_OUT=$(run_with_timeout 6 tailscale status 2>>"$LOG_FILE")
273     # Online peers = rows starting with a Tailscale IP, minus self and 'offline'.
274     # (A '-' status column means online-but-idle, so it is kept.)
275     PEERS=$(printf '%s\n' "$TS_OUT" | awk -v self="$SELF_TS_IP" \
276         '$1 ~ /^100\./ && $1 != self && $0 !~ /^([:space:]]offline(,|[:space:]]$)/ {print $1, $2}')
277     if [ -z "$PEERS" ]; then
278         warn "no online peers to ping (everything else in the tailnet is offline)"
279         record WARN "tailnet"
280     else
281         unreached=0
282         while read -r peer_ip peer_name; do
283             PONG=$(run_with_timeout 10 tailscale ping --timeout=3s --c 1 --until-direct=false "$peer_ip" 2>>"$LOG_FILE")
284             if printf '%s' "$PONG" | grep -q '~pong'; then
285                 detail=$(printf '%s\n' "$PONG" | head -1 | sed -E 's/~pong from [^ ]+ \([^\)]+\) //' )
286                 ok "$peer_name ($peer_ip) pong $detail"
287             else
288                 fail "$peer_name ($peer_ip) no pong (online per control plane, data path dead)"
289                 unreached=$((unreached+1))
290             fi
291         done <<< "$PEERS"
292         n_peers=$(printf '%s\n' "$PEERS" | wc -l | tr -d '[:space:]')
293         if [ "$unreached" -eq 0 ]; then
294             record PASS "tailnet"
295         elif [ "$unreached" -lt "$n_peers" ]; then

```



```

296     record WARN "tailnet"
297 else
298     record FAIL "tailnet"
299 fi
300 fi
301 fi
302
303 #
304 section "VERDICT"
305 #
306 pass=0; warnc=0; failc=0
307 for r in "${RESULTS[@]"; do
308     case "${r%|*}" in PASS) pass=$((pass+1));; WARN) warnc=$((warnc+1));; FAIL) failc=$((failc+1));; esac
309 done
310 for r in "${RESULTS[@]"; do
311     st="${r%|*}"; lbl="${r##*|}"
312     case "$st" in
313         PASS) ok "$lbl" ;;
314         WARN) warn "$lbl" ;;
315         FAIL) fail "$lbl" ;;
316     esac
317 done
318 emit ""
319 if [ "$failc" -eq 0 ] && [ "$warnc" -eq 0 ]; then
320     emit "${GREEN}${BOLD}SWEEP PASS gateway healthy (${pass}/${#RESULTS[@]} checks green){NC}"
321     EXIT=0
322 elif [ "$failc" -eq 0 ]; then
323     emit "${YELLOW}${BOLD}SWEEP PASS with warnings ${pass} pass, ${warnc} warn, 0 fail${NC}"
324     EXIT=0
325 else
326     emit "${RED}${BOLD}SWEEP FAIL ${pass} pass, ${warnc} warn, ${failc} fail${NC}"
327     EXIT=1
328 fi
329 emit ""
330 [ "$QUIET" = true ] || emit "Full report: $REPORT_FILE"
331 exit "$EXIT"

```

## 46 scripts/unlock-files.sh

```

1 #!/bin/bash
2 # Resolve project root relative to this script's location, so the script
3 # works regardless of the caller's CWD.
4 SCRIPT_DIR="$(cd "$(dirname "${BASH_SOURCE[0]}")" && pwd)"
5 PROJECT_ROOT="$(cd "$SCRIPT_DIR/.." && pwd)"
6
7 echo "Unlocking .env..."
8 if [ -f "$PROJECT_ROOT/.env" ]; then
9     sudo cat "$PROJECT_ROOT/.env" > "$PROJECT_ROOT/.env.tmp" && mv "$PROJECT_ROOT/.env.tmp" "$PROJECT_ROOT/.env"
10     echo ".env unlocked"
11 fi
12
13 echo "Unlocking AdGuard cache files..."
14 for f in "$PROJECT_ROOT"/adguard/work/userfilters/compiled_rules.txt \
15         "$PROJECT_ROOT"/adguard/work/userfilters/ip_blocklist.ipset \
16         "$PROJECT_ROOT"/adguard/work/userfilters/cache/*.txt \
17         "$PROJECT_ROOT"/adguard/work/userfilters/cache/ip/*.txt \
18         "$PROJECT_ROOT"/adguard/work/data/filters/*.txt; do
19     if [ -f "$f" ]; then
20         cat "$f" > "$f.tmp" && mv "$f.tmp" "$f"
21         echo "Unlocked $f"
22     fi
23 done
24
25 echo ""
26 echo "All locked files have been successfully bypassed via atomic rename!"

```

## 47 tor-bridges.txt

## 48 white\_list.txt

```
1 0.client-channel.google.com
2 1drv.com
3 2.android.pool.ntp.org
4 akamaihd.net
5 akamaitechnologies.com
6 akamaized.net
7 amazonaws.com
8 android.clients.google.com
9 api-adservices.apple.com
10 api.ipify.org
11 api.rlje.net
12 app-api.ted.com
13 appleid.apple.com
14 apps.skype.com
15 appsbackup-pa.clients6.google.com
16 appsbackup-pa.googleapis.com
17 appspot-preview.l.google.com
18 apt.sonarr.tv
19 aspnetcdn.com
20 attestation.xboxlive.com
21 ax.phobos.apple.com.edgesuite.net
22 books-analytics-events.apple.com
23 brightcove.net
24 c.s-microsoft.com
25 cdn.cloudflare.net
26 cdn.embedly.com
27 cdn.optimizely.com
28 cdn.vidible.tv
29 cdn2.optimizely.com
30 cdn3.optimizely.com
31 cdnjs.cloudflare.com
32 cert.mgt.xboxlive.com
33 clientconfig.passport.net
34 clients1.google.com
35 clients2.google.com
36 clients3.google.com
37 clients4.google.com
38 clients5.google.com
39 clients6.google.com
40 connectivitycheck.android.com
41 connectivitycheck.gstatic.com
42 continuum.dds.microsoft.com
43 cpms.spop10.ams.plex.bz
44 cpms35.spop10.ams.plex.bz
45 cse.google.com
46 ctldl.windowsupdate.com
47 cws.conviva.com
48 d2c8v52ll5s99u.cloudfront.net
49 d2gatte9o95jao.cloudfront.net
50 dashboard.plex.tv
51 dataplicity.com
52 def-vef.xboxlive.com
53 delivery.vidible.tv
54 dev.virtualearth.net
55 device.auth.xboxlive.com
56 display.ugc.bazaarvoice.com
57 displaycatalog.mp.microsoft.com
58 dl.delivery.mp.microsoft.com
59 dl.dropbox.com
60 dl.dropboxusercontent.com
61 dns.msftncsi.com
62 download.sonarr.tv
63 drift.com
64 driftt.com
65 dynupdate.no-ip.com
66 ecn.dev.virtualearth.net
67 edge.api.brightcove.com
68 eds.xboxlive.com
69 fonts.gstatic.com
70 forums.sonarr.tv
71 g.live.com
72 geo-prod.do.dsp.mp.microsoft.com
73 geo3.ggpht.com
74 gfwsl.geforce.com
75 giphy.com
76 github.com
77 github.io
78 googleapis.com
```

79 gravatar.com  
80 gstatic.com  
81 help.ui.xboxlive.com  
82 hls.ted.com  
83 i.ytimg.com  
84 il.ytimg.com  
85 iadsk.apple.com  
86 imagesak.secureserver.net  
87 img.vidible.tv  
88 imgix.net  
89 imgs.xkcd.com  
90 instantmessaging-pa.googleapis.com  
91 intercom.io  
92 jquery.com  
93 jsdelivr.net  
94 keystone.mwbsys.com  
95 lastfm-img2.akamaized.net  
96 licensing.xboxlive.com  
97 live.com  
98 livepassdl.conviva.com  
99 login.live.com  
100 login.microsoftonline.com  
101 manifest.googlevideo.com  
102 meta-db-worker02.pop.ric.plex.bz  
103 meta.plex.bz  
104 meta.plex.tv  
105 microsoftonline.com  
106 msftncsi.com  
107 my.plexapp.com  
108 nexusrules.officeapps.live.com  
109 nine.plugins.plexapp.com  
110 no-ip.com  
111 node.plexapp.com  
112 notes-analytics-events.apple.com  
113 notify.xboxlive.com  
114 npr-news.streaming.adswizz.com  
115 ns1.dropbox.com  
116 ns2.dropbox.com  
117 o1.email.plex.tv  
118 o2.sg0.plex.tv  
119 ocsp.apple.com  
120 office.com  
121 office.net  
122 office365.com  
123 officeclient.microsoft.com  
124 om.cbsi.com  
125 onedrive.live.com  
126 outlook.live.com  
127 outlook.office365.com  
128 pings.conviva.com  
129 placeholder.it  
130 placeholderit.imgix.net  
131 players.brightcove.net  
132 pricelist.skype.com  
133 products.office.com  
134 proxy.plex.bz  
135 proxy.plex.tv  
136 proxy02.pop.ord.plex.bz  
137 pubsub.plex.bz  
138 pubsub.plex.tv  
139 raw.githubusercontent.com  
140 redirector.googlevideo.com  
141 res.cloudinary.com  
142 s.gateway.messenger.live.com  
143 s.marketwatch.com  
144 s.youtube.com  
145 s.ytimg.com  
146 s1.wp.com  
147 s2.youtube.com  
148 s3.amazonaws.com  
149 sa.symcb.com  
150 secure.avangate.com  
151 secure.brightcove.com  
152 secure.surveymonkey.com  
153 services.sonarr.tv  
154 skyhook.sonarr.tv  
155 spclient.wg.spotify.com  
156 ssl.p.jwpcdn.com  
157 staging.plex.tv  
158 status.plex.tv  
159 t.co

```

160 t0.ssl.ak.dynamic.tiles.virtualearth.net
161 t0.ssl.ak.tiles.virtualearth.net
162 tawk.to
163 tedcdn.com
164 themoviedb.com
165 thetvdb.com
166 tinyurl.com
167 title.auth.xboxlive.com
168 title.mgt.xboxlive.com
169 traffic.libsyn.com
170 tvdb2.plex.tv
171 tvthemes.plexapp.com
172 twimg.com
173 ui.skype.com
174 video-stats.l.google.com
175 videos.vidible.tv
176 vidtech.cbsinteractive.com
177 weather-analytics-events.apple.com
178 widget-cdn.rpxnow.com
179 win10.ipv6.microsoft.com
180 wp.com
181 ws.audioscrobbler.com
182 www.dataplicity.com
183 www.googleapis.com
184 www.msftconnecttest.com
185 www.msftncsi.com
186 www.no-ip.com
187 www.youtube-nocookie.com
188 xbox.ipv6.microsoft.com
189 xboxexperiencesprod.experimentation.xboxlive.com
190 xflight.xboxlive.com
191 xkms.xboxlive.com
192 xsts.auth.xboxlive.com
193youtu.be
194 youtube-nocookie.com
195 yt3.ggpht.com
196 zee.cws.conviva.com
197
198 # YouTube family base domains so every subdomain (video stream servers on
199 # *.googlevideo.com, image hosts on *.ytimg.com, profile thumbnails on
200 # *.ggpht.com, etc.) is automatically allowlisted via the AdGuard '||domain^'
201 # wildcard. The specific s.youtube.com / manifest.googlevideo.com entries
202 # above are kept for clarity but will be optimized away by rule-compiler
203 # because their parents are now in this list.
204 youtube.com
205 googlevideo.com
206 ytimg.com
207 ggpht.com
208 gvt1.com
209 gvt2.com
210 gvt3.com
211 gvt1.net
212 gvt2.net
213 gvt3.net
214
215 edge-mqtt.facebook.com
216 gateway.instagram.com
217 i.instagram.com
218
219 facebook.com
220 fb.com
221 fbcdn.net
222 fbsbx.com
223 messenger.com
224 whatsapp.com
225 web.whatsapp.com
226 instagram.com
227 cdninstagram.com
228 static.whatsapp.net
229 media-bos5-1.cdn.whatsapp.net
230 media-yyz1-1.cdn.whatsapp.net
231 webtp.whatsapp.net
232 mmg.whatsapp.net
233 g.whatsapp.net
234 g-fallback.whatsapp.net

```